

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417495
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Parsons; Scott et al.

Offload processing of data packets containing financial market data

Abstract

Various techniques are disclosed for offloading the processing of data packets that contain financial market data. For example, data packets that contain financial market data can be processed through an offload processor deployed in a network element of a data distribution network to link the data packets with appropriate parsing metadata for the data packets. The network element can be an intelligent switch, such as an intelligent switch upstream or downstream from an electronic trading platform.

Inventors: Parsons; Scott (St. Charles, MO), Taylor; David E. (St. Louis, MO), Indeck; Ronald S. (St. Louis, MO)

Applicant: Exegy Incorporated (St. Louis, MO)

Family ID: 1000008820373

Assignee: Exegy Incorporated (St. Louis, MO)

Appl. No.: 17/215469

Filed: March 29, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20210217088 A1	Jul. 15, 2021

Related U.S. Application Data

continuation parent-doc US 16180377 20181105 US 10963962 child-doc US 17215469
division parent-doc US 13833098 20130315 US 10121196 20181106 child-doc US 16180377
us-provisional-application US 61616181 20120327

Publication Classification

Int. Cl.: G06Q40/04 (20120101)

U.S. Cl.:

Field of Classification Search

USPC: None

References Cited**U.S. PATENT DOCUMENTS**

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2046381	12/1935	Hicks et al.	N/A	N/A
3082402	12/1962	Scantlin	N/A	N/A
3296597	12/1966	Scantlin et al.	N/A	N/A
3573747	12/1970	Adams et al.	N/A	N/A
3581072	12/1970	Nymeyer	N/A	N/A
3601808	12/1970	Vlack	N/A	N/A
3611314	12/1970	Pritchard, Jr. et al.	N/A	N/A
3729712	12/1972	Glassman	N/A	N/A
3824375	12/1973	Gross et al.	N/A	N/A
3848235	12/1973	Lewis et al.	N/A	N/A
3906455	12/1974	Houston et al.	N/A	N/A
4044334	12/1976	Bachman et al.	N/A	N/A
4081607	12/1977	Vitols et al.	N/A	N/A
4298898	12/1980	Cardot	N/A	N/A
4300193	12/1980	Bradley et al.	N/A	N/A
4314356	12/1981	Scarbrough	N/A	N/A
4385393	12/1982	Chaure et al.	N/A	N/A
4412287	12/1982	Braddock, III	N/A	N/A
4464718	12/1983	Dixon et al.	N/A	N/A
4550436	12/1984	Freeman et al.	N/A	N/A
4674044	12/1986	Kalmus et al.	N/A	N/A
4811214	12/1988	Nosenchuck et al.	N/A	N/A
4823306	12/1988	Barbic et al.	N/A	N/A
4868866	12/1988	Williams, Jr.	N/A	N/A
4903201	12/1989	Wagner	N/A	N/A
4941178	12/1989	Chuang	N/A	N/A
5023910	12/1990	Thomson	N/A	N/A
5038284	12/1990	Kramer	N/A	N/A
5050075	12/1990	Herman et al.	N/A	N/A
5063507	12/1990	Lindsey et al.	N/A	N/A
5077665	12/1990	Silverman et al.	N/A	N/A
5101353	12/1991	Lupien et al.	N/A	N/A
5101424	12/1991	Clayton et al.	N/A	N/A
5126936	12/1991	Champion et al.	N/A	N/A
5140692	12/1991	Morita	N/A	N/A
5161103	12/1991	Kosaka et al.	N/A	N/A
5163131	12/1991	Row et al.	N/A	N/A
5179626	12/1992	Thomson	N/A	N/A
5226165	12/1992	Martin	N/A	N/A
5243655	12/1992	Wang	N/A	N/A
5249292	12/1992	Chiappa	N/A	N/A
5255136	12/1992	Machado et al.	N/A	N/A

5258908	12/1992	Hartheimer et al.	N/A	N/A
5265065	12/1992	Turtle	N/A	N/A
5267148	12/1992	Kosaka et al.	N/A	N/A
5270922	12/1992	Higgins	N/A	N/A
5297032	12/1993	Trojan et al.	N/A	N/A
5313560	12/1993	Maruoka et al.	N/A	N/A
5315634	12/1993	Tanaka et al.	N/A	N/A
5319776	12/1993	Hile et al.	N/A	N/A
5327521	12/1993	Savic et al.	N/A	N/A
5339411	12/1993	Heaton, Jr.	N/A	N/A
5347634	12/1993	Herrell et al.	N/A	N/A
5371794	12/1993	Diffie et al.	N/A	N/A
5375055	12/1993	Togher et al.	N/A	N/A
5388259	12/1994	Fleischman et al.	N/A	N/A
5396253	12/1994	Chia	N/A	N/A
5404488	12/1994	Kerrigan et al.	N/A	N/A
5418951	12/1994	Damashek	N/A	N/A
5421028	12/1994	Swanson	N/A	N/A
5432822	12/1994	Kaewell, Jr.	N/A	N/A
5461712	12/1994	Chelstowski et al.	N/A	N/A
5465353	12/1994	Hull et al.	N/A	N/A
5481735	12/1995	Mortensen et al.	N/A	N/A
5488725	12/1995	Turtle et al.	N/A	N/A
5497317	12/1995	Hawkins et al.	N/A	N/A
5497488	12/1995	Akizawa et al.	N/A	N/A
5500793	12/1995	Deming, Jr. et al.	N/A	N/A
5517642	12/1995	Bezek et al.	N/A	N/A
5544352	12/1995	Egger	N/A	N/A
5546578	12/1995	Takada et al.	N/A	N/A
5596569	12/1996	Madonna et al.	N/A	N/A
5619574	12/1996	Johnson et al.	N/A	N/A
5651125	12/1996	Witt et al.	N/A	N/A
5680634	12/1996	Estes	N/A	N/A
5684980	12/1996	Casselman	N/A	N/A
5701464	12/1996	Aucsmith	N/A	N/A
5712942	12/1997	Jennings et al.	N/A	N/A
5721898	12/1997	Beardsley et al.	N/A	N/A
5740244	12/1997	Indeck et al.	N/A	N/A
5740466	12/1997	Geldman et al.	N/A	N/A
5774835	12/1997	Ozawa et al.	N/A	N/A
5774839	12/1997	Shlomot	N/A	N/A
5781772	12/1997	Wilkinson, III et al.	N/A	N/A
5781921	12/1997	Nichols	N/A	N/A
5802290	12/1997	Casselman	N/A	N/A
5805832	12/1997	Brown et al.	N/A	N/A
5809483	12/1997	Broka et al.	N/A	N/A
5813000	12/1997	Furlani	N/A	N/A
5819273	12/1997	Vora et al.	N/A	N/A
5819290	12/1997	Fujita et al.	N/A	N/A
5826075	12/1997	Bealkowski et al.	N/A	N/A
5845266	12/1997	Lupien et al.	N/A	N/A
5857176	12/1998	Ginsberg	N/A	N/A
5864738	12/1998	Kessler et al.	N/A	N/A

5870730	12/1998	Furuya et al.	N/A	N/A
5873071	12/1998	Ferstenberg et al.	N/A	N/A
5884286	12/1998	Daughtery, III	N/A	N/A
5905974	12/1998	Fraser et al.	N/A	N/A
5913211	12/1998	Nitta	N/A	N/A
5930753	12/1998	Potamianos et al.	N/A	N/A
5943421	12/1998	Grabon	N/A	N/A
5943429	12/1998	Händel	N/A	N/A
5950006	12/1998	Crater et al.	N/A	N/A
5963923	12/1998	Garber	N/A	N/A
5978801	12/1998	Yuasa	N/A	N/A
5987432	12/1998	Zusman et al.	N/A	N/A
5991881	12/1998	Conklin et al.	N/A	N/A
5995963	12/1998	Nanba et al.	N/A	N/A
6006264	12/1998	Colby et al.	N/A	N/A
6016483	12/1999	Rickard et al.	N/A	N/A
6023755	12/1999	Casselman	N/A	N/A
6023760	12/1999	Karttunen	N/A	N/A
6028939	12/1999	Yin	N/A	N/A
6034538	12/1999	Abramovici	N/A	N/A
6044407	12/1999	Jones et al.	N/A	N/A
6058391	12/1999	Gardner	N/A	N/A
6061662	12/1999	Makivic	N/A	N/A
6064739	12/1999	Davis	N/A	N/A
6067569	12/1999	Khaki et al.	N/A	N/A
6070172	12/1999	Lowe	N/A	N/A
6073160	12/1999	Grantham et al.	N/A	N/A
6084584	12/1999	Nahi et al.	N/A	N/A
6096091	12/1999	Hartmann	N/A	N/A
6105067	12/1999	Batra	N/A	N/A
6134551	12/1999	Aucsmith	N/A	N/A
6138176	12/1999	McDonald et al.	N/A	N/A
6147976	12/1999	Shand et al.	N/A	N/A
6169969	12/2000	Cohen	N/A	N/A
6173270	12/2000	Cristofich et al.	N/A	N/A
6173276	12/2000	Kant et al.	N/A	N/A
6178494	12/2000	Casselman	N/A	N/A
6195024	12/2000	Fallon	N/A	N/A
6226676	12/2000	Crump et al.	N/A	N/A
6236980	12/2000	Reese	N/A	N/A
6243753	12/2000	Machin et al.	N/A	N/A
6247060	12/2000	Boucher et al.	N/A	N/A
6263321	12/2000	Daughtery, III	N/A	N/A
6272616	12/2000	Fernando et al.	N/A	N/A
6278982	12/2000	Korhammer et al.	N/A	N/A
6279113	12/2000	Vaidya	N/A	N/A
6279140	12/2000	Slane	N/A	N/A
6289440	12/2000	Casselman	N/A	N/A
6295530	12/2000	Ritchie et al.	N/A	N/A
6304858	12/2000	Mosler et al.	N/A	N/A
6309424	12/2000	Fallon	N/A	N/A
6317728	12/2000	Kane	N/A	N/A
6317795	12/2000	Malkin et al.	N/A	N/A

6321258	12/2000	Stollfus et al.	N/A	N/A
6336150	12/2001	Ellis et al.	N/A	N/A
6339819	12/2001	Huppenthal et al.	N/A	N/A
6370592	12/2001	Kumpf	N/A	N/A
6370645	12/2001	Lee et al.	N/A	N/A
6377942	12/2001	Hinsley et al.	N/A	N/A
6388584	12/2001	Dorward et al.	N/A	N/A
6397259	12/2001	Lincke et al.	N/A	N/A
6397335	12/2001	Franczek et al.	N/A	N/A
6412000	12/2001	Riddle et al.	N/A	N/A
6415269	12/2001	Dinwoodie	N/A	N/A
6418419	12/2001	Nieboer et al.	N/A	N/A
6430272	12/2001	Maruyama et al.	N/A	N/A
6456982	12/2001	Pilipovic	N/A	N/A
6463474	12/2001	Fuh et al.	N/A	N/A
6484209	12/2001	Momirov	N/A	N/A
6499107	12/2001	Glechauf et al.	N/A	N/A
6535868	12/2002	Galeazzi et al.	N/A	N/A
6546375	12/2002	Pang et al.	N/A	N/A
6578147	12/2002	Shanklin et al.	N/A	N/A
6581098	12/2002	Kumpf	N/A	N/A
6591302	12/2002	Boucher et al.	N/A	N/A
6594643	12/2002	Freeny, Jr.	N/A	N/A
6597812	12/2002	Fallon et al.	N/A	N/A
6601094	12/2002	Mentze et al.	N/A	N/A
6601104	12/2002	Fallon	N/A	N/A
6604158	12/2002	Fallon	N/A	N/A
6624761	12/2002	Fallon	N/A	N/A
6625150	12/2002	Yu	N/A	N/A
6691301	12/2003	Bowen	N/A	N/A
6704816	12/2003	Burke	N/A	N/A
6710702	12/2003	Averbuch et al.	N/A	N/A
6711558	12/2003	Indeck et al.	N/A	N/A
6765918	12/2003	Dixon et al.	N/A	N/A
6766304	12/2003	Kemp, II et al.	N/A	N/A
6772132	12/2003	Kemp, II et al.	N/A	N/A
6772136	12/2003	Kant et al.	N/A	N/A
6772345	12/2003	Shetty	N/A	N/A
6778968	12/2003	Gulati	N/A	N/A
6785677	12/2003	Fritchman	N/A	N/A
6804667	12/2003	Martin	N/A	N/A
6807156	12/2003	Veres et al.	N/A	N/A
6820129	12/2003	Courey, Jr.	N/A	N/A
6839686	12/2004	Galant	N/A	N/A
6847645	12/2004	Potter et al.	N/A	N/A
6850906	12/2004	Chadha et al.	N/A	N/A
6870837	12/2004	Ho et al.	N/A	N/A
6877044	12/2004	Lo et al.	N/A	N/A
6886103	12/2004	Brustoloni et al.	N/A	N/A
6901461	12/2004	Bennett	N/A	N/A
6931408	12/2004	Adams et al.	N/A	N/A
6931545	12/2004	Ta et al.	N/A	N/A
6944168	12/2004	Paatela et al.	N/A	N/A

6978223	12/2004	Milliken	N/A	N/A
6981054	12/2004	Krishna	N/A	N/A
7003488	12/2005	Dunne et al.	N/A	N/A
7024384	12/2005	Daughtery, III	N/A	N/A
7046848	12/2005	Olcott	N/A	N/A
7058735	12/2005	Spencer	N/A	N/A
7065475	12/2005	Brundobler	N/A	N/A
7065482	12/2005	Shorey et al.	N/A	N/A
7089206	12/2005	Martin	N/A	N/A
7089326	12/2005	Boucher et al.	N/A	N/A
7093023	12/2005	Lockwood et al.	N/A	N/A
7099838	12/2005	Gastineau et al.	N/A	N/A
7101188	12/2005	Summers et al.	N/A	N/A
7103569	12/2005	Groveman et al.	N/A	N/A
7117280	12/2005	Vasudevan	N/A	N/A
7124106	12/2005	Stallaert et al.	N/A	N/A
7127424	12/2005	Kemp, II et al.	N/A	N/A
7130913	12/2005	Fallon	N/A	N/A
7139743	12/2005	Indeck et al.	N/A	N/A
7149715	12/2005	Browne et al.	N/A	N/A
7161506	12/2006	Fallon	N/A	N/A
7167980	12/2006	Chiu	N/A	N/A
7177833	12/2006	Marynowski et al.	N/A	N/A
7181437	12/2006	Indeck et al.	N/A	N/A
7181608	12/2006	Fallon et al.	N/A	N/A
7191233	12/2006	Miller	N/A	N/A
7212998	12/2006	Muller et al.	N/A	N/A
7219125	12/2006	Day	N/A	N/A
7222114	12/2006	Chan et al.	N/A	N/A
7224185	12/2006	Campbell et al.	N/A	N/A
7225188	12/2006	Gai et al.	N/A	N/A
7228289	12/2006	Brumfield et al.	N/A	N/A
7249118	12/2006	Sandler et al.	N/A	N/A
7251629	12/2006	Marynowski et al.	N/A	N/A
7257842	12/2006	Barton et al.	N/A	N/A
7277887	12/2006	Burrows et al.	N/A	N/A
7287037	12/2006	An et al.	N/A	N/A
7305383	12/2006	Kubesh et al.	N/A	N/A
7305391	12/2006	Wyschogrod et al.	N/A	N/A
7321937	12/2007	Fallon	N/A	N/A
7356498	12/2007	Kaminsky et al.	N/A	N/A
7363277	12/2007	Dutta et al.	N/A	N/A
7372875	12/2007	Hadzic et al.	N/A	N/A
7376755	12/2007	Pandya	N/A	N/A
7378992	12/2007	Fallon	N/A	N/A
7386046	12/2007	Fallon et al.	N/A	N/A
7406444	12/2007	Eng et al.	N/A	N/A
7415723	12/2007	Pandya	N/A	N/A
7417568	12/2007	Fallon et al.	N/A	N/A
7420931	12/2007	Nanda et al.	N/A	N/A
7447788	12/2007	Ahmed et al.	N/A	N/A
7454418	12/2007	Wang	N/A	N/A
7457834	12/2007	Jung et al.	N/A	N/A

7461064	12/2007	Fontoura et al.	N/A	N/A
7472237	12/2007	Herbst et al.	N/A	N/A
7478431	12/2008	Nachenberg	N/A	N/A
7480253	12/2008	Allan	N/A	N/A
7487264	12/2008	Pandya	N/A	N/A
7487327	12/2008	Chang et al.	N/A	N/A
7496108	12/2008	Biran et al.	N/A	N/A
7536462	12/2008	Pandya	N/A	N/A
7539845	12/2008	Wentzlaff et al.	N/A	N/A
7558753	12/2008	Neubert et al.	N/A	N/A
7558925	12/2008	Bouchard et al.	N/A	N/A
7565525	12/2008	Vorbach et al.	N/A	N/A
7580719	12/2008	Karmarkar	N/A	N/A
7587476	12/2008	Sato	N/A	N/A
7603303	12/2008	Kraus et al.	N/A	N/A
7606267	12/2008	Ho et al.	N/A	N/A
7606968	12/2008	Branscome et al.	N/A	N/A
7617291	12/2008	Fan et al.	N/A	N/A
7627693	12/2008	Pandya	N/A	N/A
7631107	12/2008	Pandya	N/A	N/A
7636703	12/2008	Taylor	N/A	N/A
7660761	12/2009	Zhou et al.	N/A	N/A
7668849	12/2009	Narancic et al.	N/A	N/A
7685121	12/2009	Brown et al.	N/A	N/A
7685254	12/2009	Pandya	N/A	N/A
7698338	12/2009	Hinshaw et al.	N/A	N/A
7701945	12/2009	Roesch et al.	N/A	N/A
7714747	12/2009	Fallon	N/A	N/A
7715436	12/2009	Eiriksson et al.	N/A	N/A
7760733	12/2009	Eiriksson et al.	N/A	N/A
7761459	12/2009	Zhang et al.	N/A	N/A
7788293	12/2009	Pasztor et al.	N/A	N/A
7827190	12/2009	Pandya	N/A	N/A
7831606	12/2009	Pandya	N/A	N/A
7831607	12/2009	Pandya	N/A	N/A
7831720	12/2009	Noureddine et al.	N/A	N/A
7840482	12/2009	Singla et al.	N/A	N/A
7856545	12/2009	Casselman	N/A	N/A
7856546	12/2009	Casselman et al.	N/A	N/A
7869442	12/2010	Kamboh et al.	N/A	N/A
7870217	12/2010	Pandya	N/A	N/A
7890692	12/2010	Pandya	N/A	N/A
7899976	12/2010	Pandya	N/A	N/A
7899977	12/2010	Pandya	N/A	N/A
7899978	12/2010	Pandya	N/A	N/A
7908213	12/2010	Monroe et al.	N/A	N/A
7908259	12/2010	Branscome et al.	N/A	N/A
7912808	12/2010	Pandya	N/A	N/A
7913022	12/2010	Baxter	N/A	N/A
7917299	12/2010	Buhler et al.	N/A	N/A
7921046	12/2010	Parsons et al.	N/A	N/A
7944920	12/2010	Pandya	N/A	N/A
7945528	12/2010	Cytron et al.	N/A	N/A

7949650	12/2010	Indeck et al.	N/A	N/A
7953743	12/2010	Indeck et al.	N/A	N/A
7954114	12/2010	Chamberlain et al.	N/A	N/A
7991667	12/2010	Kraus et al.	N/A	N/A
7996348	12/2010	Pandya	N/A	N/A
8005966	12/2010	Pandya	N/A	N/A
8015099	12/2010	Reid	N/A	N/A
8024253	12/2010	Peterffy et al.	N/A	N/A
8027893	12/2010	Burrows et al.	N/A	N/A
8032440	12/2010	Hait	N/A	N/A
8046283	12/2010	Burns et al.	N/A	N/A
8051022	12/2010	Pandya	N/A	N/A
8055601	12/2010	Pandya	N/A	N/A
8069102	12/2010	Indeck et al.	N/A	N/A
8073763	12/2010	Merrin et al.	N/A	N/A
8095508	12/2011	Chamberlain et al.	N/A	N/A
8131697	12/2011	Indeck et al.	N/A	N/A
8140416	12/2011	Borkovec et al.	N/A	N/A
8156101	12/2011	Indeck et al.	N/A	N/A
8181239	12/2011	Pandya	N/A	N/A
8200599	12/2011	Pandya	N/A	N/A
8224800	12/2011	Branscome et al.	N/A	N/A
8229918	12/2011	Branscome et al.	N/A	N/A
8234267	12/2011	Branscome et al.	N/A	N/A
8244718	12/2011	Chamdani et al.	N/A	N/A
8260764	12/2011	Gruber	N/A	N/A
8281026	12/2011	Lankford et al.	N/A	N/A
8326819	12/2011	Indeck et al.	N/A	N/A
8374986	12/2012	Indeck et al.	N/A	N/A
8407122	12/2012	Parsons et al.	N/A	N/A
8458081	12/2012	Parsons et al.	N/A	N/A
8478680	12/2012	Parsons et al.	N/A	N/A
8479124	12/2012	Baxter et al.	N/A	N/A
8515682	12/2012	Buhler et al.	N/A	N/A
8549024	12/2012	Indeck et al.	N/A	N/A
8595104	12/2012	Parsons et al.	N/A	N/A
8600856	12/2012	Parsons et al.	N/A	N/A
8601086	12/2012	Pandya	N/A	N/A
8620881	12/2012	Chamberlain et al.	N/A	N/A
8626624	12/2013	Parsons et al.	N/A	N/A
8655764	12/2013	Parsons et al.	N/A	N/A
8660925	12/2013	Borkovec et al.	N/A	N/A
8737606	12/2013	Taylor et al.	N/A	N/A
8751452	12/2013	Chamberlain et al.	N/A	N/A
8762249	12/2013	Taylor et al.	N/A	N/A
8768805	12/2013	Taylor et al.	N/A	N/A
8768888	12/2013	Chamberlain et al.	N/A	N/A
8843408	12/2013	Singla et al.	N/A	N/A
8879727	12/2013	Taylor et al.	N/A	N/A
8880501	12/2013	Indeck et al.	N/A	N/A
8880551	12/2013	Hinshaw et al.	N/A	N/A
9020928	12/2014	Indeck et al.	N/A	N/A
9047243	12/2014	Taylor et al.	N/A	N/A

9166597	12/2014	Denisenko et al.	N/A	N/A
9176775	12/2014	Chamberlain et al.	N/A	N/A
9378161	12/2015	Dalal et al.	N/A	N/A
9961006	12/2017	Sutardja et al.	N/A	N/A
10037568	12/2017	Taylor et al.	N/A	N/A
10121196	12/2017	Parsons et al.	N/A	N/A
10191974	12/2018	Indeck et al.	N/A	N/A
10229453	12/2018	Taylor et al.	N/A	N/A
10572824	12/2019	Chamberlain et al.	N/A	N/A
10650452	12/2019	Parsons et al.	N/A	N/A
10909623	12/2020	Indeck et al.	N/A	N/A
10929152	12/2020	Chamberlain et al.	N/A	N/A
10957423	12/2020	Buhler et al.	N/A	N/A
11275594	12/2021	Chamberlain et al.	N/A	N/A
11416778	12/2021	Chamberlain et al.	N/A	N/A
11436672	12/2021	Parsons et al.	N/A	N/A
11449538	12/2021	Indeck et al.	N/A	N/A
11803912	12/2022	Taylor et al.	N/A	N/A
12148032	12/2023	Parsons et al.	N/A	N/A
2001/0003193	12/2000	Woodring et al.	N/A	N/A
2001/0004354	12/2000	Jolitz	N/A	N/A
2001/0013048	12/2000	Imbert de Tremiolles et al.	N/A	N/A
2001/0025315	12/2000	Jolitz	N/A	N/A
2001/0042040	12/2000	Keith	N/A	N/A
2001/0044770	12/2000	Keith	N/A	N/A
2001/0047473	12/2000	Fallon	N/A	N/A
2001/0052038	12/2000	Fallon et al.	N/A	N/A
2001/0056547	12/2000	Dixon	N/A	N/A
2002/0004820	12/2001	Baldwin et al.	N/A	N/A
2002/0010825	12/2001	Wilson	N/A	N/A
2002/0019812	12/2001	Board et al.	N/A	N/A
2002/0023010	12/2001	Rittmaster et al.	N/A	N/A
2002/0031125	12/2001	Sato	N/A	N/A
2002/0038276	12/2001	Buhannic et al.	N/A	N/A
2002/0049841	12/2001	Johnson et al.	N/A	N/A
2002/0054604	12/2001	Kadambi et al.	N/A	N/A
2002/0069370	12/2001	Mack	N/A	N/A
2002/0069375	12/2001	Bowen	N/A	N/A
2002/0072893	12/2001	Wilson	N/A	N/A
2002/0080871	12/2001	Fallon et al.	N/A	N/A
2002/0082967	12/2001	Kaminsky et al.	N/A	N/A
2002/0091826	12/2001	Comeau et al.	N/A	N/A
2002/0095512	12/2001	Rana et al.	N/A	N/A
2002/0095519	12/2001	Philbrick et al.	N/A	N/A
2002/0100029	12/2001	Bowen	N/A	N/A
2002/0101425	12/2001	Hamid	N/A	N/A
2002/0105911	12/2001	Pruthi et al.	N/A	N/A
2002/0129140	12/2001	Peled et al.	N/A	N/A
2002/0138376	12/2001	Hinkle	N/A	N/A
2002/0143521	12/2001	Call	N/A	N/A
2002/0150248	12/2001	Kovacevic	N/A	N/A
2002/0156998	12/2001	Casselman	N/A	N/A
2002/0162025	12/2001	Sutton et al.	N/A	N/A

2002/0166063	12/2001	Lachman et al.	N/A	N/A
2002/0169873	12/2001	Zodnik	N/A	N/A
2002/0180742	12/2001	Hamid	N/A	N/A
2002/0198813	12/2001	Patterson et al.	N/A	N/A
2002/0199173	12/2001	Bowen	N/A	N/A
2003/0009411	12/2002	Ram et al.	N/A	N/A
2003/0009693	12/2002	Brock et al.	N/A	N/A
2003/0014521	12/2002	Elson et al.	N/A	N/A
2003/0014662	12/2002	Gupta et al.	N/A	N/A
2003/0018630	12/2002	Indeck et al.	N/A	N/A
2003/0023653	12/2002	Dunlop et al.	N/A	N/A
2003/0023876	12/2002	Bardsley et al.	N/A	N/A
2003/0028408	12/2002	RuDusky	N/A	N/A
2003/0028690	12/2002	Appleby-Alis et al.	N/A	N/A
2003/0028864	12/2002	Bowen	N/A	N/A
2003/0033234	12/2002	RuDusky	N/A	N/A
2003/0033240	12/2002	Balson et al.	N/A	N/A
2003/0033450	12/2002	Appleby-Alis	N/A	N/A
2003/0033514	12/2002	Appleby-Allis et al.	N/A	N/A
2003/0033588	12/2002	Alexander	N/A	N/A
2003/0033594	12/2002	Bowen	N/A	N/A
2003/0035547	12/2002	Newton	N/A	N/A
2003/0037037	12/2002	Adams et al.	N/A	N/A
2003/0037321	12/2002	Bowen	N/A	N/A
2003/0041129	12/2002	Applcby-Allis	N/A	N/A
2003/0043805	12/2002	Graham et al.	N/A	N/A
2003/0046668	12/2002	Bowen	N/A	N/A
2003/0051043	12/2002	Wyschogrod et al.	N/A	N/A
2003/0055658	12/2002	RuDusky	N/A	N/A
2003/0055769	12/2002	RuDusky	N/A	N/A
2003/0055770	12/2002	RuDusky	N/A	N/A
2003/0055771	12/2002	RuDusky	N/A	N/A
2003/0055777	12/2002	Ginsberg	N/A	N/A
2003/0061409	12/2002	RuDusky	N/A	N/A
2003/0065607	12/2002	Satchwell	N/A	N/A
2003/0065943	12/2002	Geis et al.	N/A	N/A
2003/0074177	12/2002	Bowen	N/A	N/A
2003/0074582	12/2002	Patel et al.	N/A	N/A
2003/0078865	12/2002	Lee	N/A	N/A
2003/0079060	12/2002	Dunlop	N/A	N/A
2003/0086300	12/2002	Noyes et al.	N/A	N/A
2003/0093343	12/2002	Huttenlocher et al.	N/A	N/A
2003/0093347	12/2002	Gray	N/A	N/A
2003/0097481	12/2002	Richter	N/A	N/A
2003/0099254	12/2002	Richter	N/A	N/A
2003/0105620	12/2002	Bowen	N/A	N/A
2003/0105721	12/2002	Ginter et al.	N/A	N/A
2003/0110229	12/2002	Kulig et al.	N/A	N/A
2003/0115485	12/2002	Milliken	N/A	N/A
2003/0117971	12/2002	Aubury	N/A	N/A
2003/0120460	12/2002	Aubury	N/A	N/A
2003/0121010	12/2002	Aubury	N/A	N/A
2003/0126065	12/2002	Eng et al.	N/A	N/A

2003/0140337	12/2002	Aubury	N/A	N/A
2003/0154284	12/2002	Bernardin et al.	N/A	N/A
2003/0154368	12/2002	Stevens et al.	N/A	N/A
2003/0163715	12/2002	Wong	N/A	N/A
2003/0172017	12/2002	Feingold et al.	N/A	N/A
2003/0177253	12/2002	Schuehler et al.	N/A	N/A
2003/0184593	12/2002	Dunlop	N/A	N/A
2003/0187662	12/2002	Wilson	N/A	N/A
2003/0191876	12/2002	Fallon	N/A	N/A
2003/0208430	12/2002	Gershon	N/A	N/A
2003/0221013	12/2002	Lockwood et al.	N/A	N/A
2003/0233302	12/2002	Weber et al.	N/A	N/A
2004/0000928	12/2003	Cheng et al.	N/A	N/A
2004/0010612	12/2003	Pandya	N/A	N/A
2004/0015502	12/2003	Alexander et al.	N/A	N/A
2004/0015633	12/2003	Smith	N/A	N/A
2004/0019703	12/2003	Burton	N/A	N/A
2004/0028047	12/2003	Hou et al.	N/A	N/A
2004/0034587	12/2003	Amberson et al.	N/A	N/A
2004/0049596	12/2003	Schuehler et al.	N/A	N/A
2004/0054924	12/2003	Chuah et al.	N/A	N/A
2004/0059666	12/2003	Waelbroeck et al.	N/A	N/A
2004/0062245	12/2003	Sharp et al.	N/A	N/A
2004/0064737	12/2003	Milliken et al.	N/A	N/A
2004/0073703	12/2003	Boucher et al.	N/A	N/A
2004/0105458	12/2003	Ishizuka	N/A	N/A
2004/0111632	12/2003	Halperin	N/A	N/A
2004/0162826	12/2003	Wyschogrod et al.	N/A	N/A
2004/0170070	12/2003	Rapp et al.	N/A	N/A
2004/0177340	12/2003	Hsu et al.	N/A	N/A
2004/0186804	12/2003	Chakraborty et al.	N/A	N/A
2004/0186814	12/2003	Chalermkraivuth et al.	N/A	N/A
2004/0190526	12/2003	Kumar et al.	N/A	N/A
2004/0199448	12/2003	Chalermkraivuth et al.	N/A	N/A
2004/0199452	12/2003	Johnston et al.	N/A	N/A
2004/0205149	12/2003	Dillon et al.	N/A	N/A
2005/0005145	12/2004	Teixeira	N/A	N/A
2005/0027634	12/2004	Gershon	N/A	N/A
2005/0033672	12/2004	Lasry et al.	N/A	N/A
2005/0038946	12/2004	Borden	N/A	N/A
2005/0044344	12/2004	Stevens	N/A	N/A
2005/0074033	12/2004	Chauveau	N/A	N/A
2005/0080649	12/2004	Alvarez et al.	N/A	N/A
2005/0086520	12/2004	Dharmapurikar et al.	N/A	N/A
2005/0091142	12/2004	Renton et al.	N/A	N/A
2005/0097027	12/2004	Kavanaugh	N/A	N/A
2005/0108518	12/2004	Pandya	N/A	N/A
2005/0111363	12/2004	Snelgrove	370/230	H04M 15/00
2005/0125551	12/2004	Oh et al.	N/A	N/A
2005/0131790	12/2004	Benzschawel et al.	N/A	N/A
2005/0135608	12/2004	Zheng	N/A	N/A
2005/0187844	12/2004	Chalermkraivuth et al.	N/A	N/A
2005/0187845	12/2004	Eklund et al.	N/A	N/A

2005/0187846	12/2004	Subbu et al.	N/A	N/A
2005/0187847	12/2004	Bonissone et al.	N/A	N/A
2005/0187848	12/2004	Bonissone et al.	N/A	N/A
2005/0187849	12/2004	Bollapragada et al.	N/A	N/A
2005/0190787	12/2004	Kuik et al.	N/A	N/A
2005/0195832	12/2004	Dharmapurikar et al.	N/A	N/A
2005/0197916	12/2004	Newell	715/225	G06Q 40/04
2005/0197938	12/2004	Davie et al.	N/A	N/A
2005/0197939	12/2004	Davie et al.	N/A	N/A
2005/0197948	12/2004	Davie et al.	N/A	N/A
2005/0216384	12/2004	Partlow et al.	N/A	N/A
2005/0228735	12/2004	Duquette	N/A	N/A
2005/0229254	12/2004	Singh et al.	N/A	N/A
2005/0240510	12/2004	Schweickert et al.	N/A	N/A
2005/0243824	12/2004	Abbazia et al.	N/A	N/A
2005/0267836	12/2004	Crosthwaite et al.	N/A	N/A
2005/0283423	12/2004	Moser et al.	N/A	N/A
2005/0283743	12/2004	Mulholland et al.	N/A	N/A
2006/0020536	12/2005	Renton et al.	N/A	N/A
2006/0020715	12/2005	Jungck	N/A	N/A
2006/0026090	12/2005	Balabon	N/A	N/A
2006/0031154	12/2005	Noviello et al.	N/A	N/A
2006/0031156	12/2005	Noviello et al.	N/A	N/A
2006/0039287	12/2005	Hasegawa et al.	N/A	N/A
2006/0047636	12/2005	Mohania et al.	N/A	N/A
2006/0053295	12/2005	Madhusudan et al.	N/A	N/A
2006/0059064	12/2005	Glinberg et al.	N/A	N/A
2006/0059065	12/2005	Glinberg et al.	N/A	N/A
2006/0059066	12/2005	Glinberg et al.	N/A	N/A
2006/0059067	12/2005	Glinberg et al.	N/A	N/A
2006/0059068	12/2005	Glinberg et al.	N/A	N/A
2006/0059069	12/2005	Glinberg et al.	N/A	N/A
2006/0059083	12/2005	Friesen et al.	N/A	N/A
2006/0109798	12/2005	Yamada	N/A	N/A
2006/0123425	12/2005	Ramarao et al.	N/A	N/A
2006/0129745	12/2005	Thiel et al.	N/A	N/A
2006/0143099	12/2005	Partlow et al.	N/A	N/A
2006/0146991	12/2005	Thompson et al.	N/A	N/A
2006/0215691	12/2005	Kobayashi et al.	N/A	N/A
2006/0242123	12/2005	Williams, Jr.	N/A	N/A
2006/0259407	12/2005	Rosenthal et al.	N/A	N/A
2006/0259417	12/2005	Marynowski et al.	N/A	N/A
2006/0269148	12/2005	Farber et al.	N/A	N/A
2006/0282281	12/2005	Egetoft	N/A	N/A
2006/0282369	12/2005	White	N/A	N/A
2006/0294059	12/2005	Chamberlain et al.	N/A	N/A
2007/0011183	12/2006	Langseth et al.	N/A	N/A
2007/0011317	12/2006	Brandyburg et al.	N/A	N/A
2007/0011687	12/2006	Ilik et al.	N/A	N/A
2007/0025351	12/2006	Cohen	N/A	N/A
2007/0061231	12/2006	Kim-E	N/A	N/A
2007/0061241	12/2006	Jovanovic et al.	N/A	N/A
2007/0061594	12/2006	Ginter et al.	N/A	N/A

2007/0067108	12/2006	Buhler et al.	N/A	N/A
2007/0067481	12/2006	Sharma et al.	N/A	N/A
2007/0078837	12/2006	Indeck et al.	N/A	N/A
2007/0094199	12/2006	Deshpande et al.	N/A	N/A
2007/0112837	12/2006	Houh et al.	N/A	N/A
2007/0115986	12/2006	Shankara	N/A	N/A
2007/0118457	12/2006	Peterffy et al.	N/A	N/A
2007/0118500	12/2006	Indeck et al.	N/A	N/A
2007/0130140	12/2006	Cytron et al.	N/A	N/A
2007/0156574	12/2006	Marynowski et al.	N/A	N/A
2007/0174841	12/2006	Chamberlain et al.	N/A	N/A
2007/0179935	12/2006	Lee et al.	N/A	N/A
2007/0198523	12/2006	Hayim	N/A	N/A
2007/0209068	12/2006	Ansari et al.	N/A	N/A
2007/0237327	12/2006	Taylor et al.	N/A	N/A
2007/0244859	12/2006	Trippe et al.	N/A	N/A
2007/0260602	12/2006	Taylor	N/A	N/A
2007/0260814	12/2006	Branscome et al.	N/A	N/A
2007/0277036	12/2006	Chamberlain et al.	N/A	N/A
2007/0294157	12/2006	Singla et al.	N/A	N/A
2007/0294162	12/2006	Borkovec	N/A	N/A
2008/0077793	12/2007	Tan et al.	N/A	N/A
2008/0082502	12/2007	Gupta	N/A	N/A
2008/0084573	12/2007	Horowitz et al.	N/A	N/A
2008/0086274	12/2007	Chamberlain et al.	N/A	N/A
2008/0097893	12/2007	Walsky et al.	N/A	N/A
2008/0104542	12/2007	Cohen et al.	N/A	N/A
2008/0109413	12/2007	Indeck et al.	N/A	N/A
2008/0114724	12/2007	Indeck et al.	N/A	N/A
2008/0114725	12/2007	Indeck et al.	N/A	N/A
2008/0114760	12/2007	Indeck et al.	N/A	N/A
2008/0126320	12/2007	Indeck et al.	N/A	N/A
2008/0133453	12/2007	Indeck et al.	N/A	N/A
2008/0133519	12/2007	Indeck et al.	N/A	N/A
2008/0162378	12/2007	Levine et al.	N/A	N/A
2008/0175239	12/2007	Sistanizadeh et al.	N/A	N/A
2008/0183688	12/2007	Chamdani et al.	N/A	N/A
2008/0189251	12/2007	Branscome et al.	N/A	N/A
2008/0189252	12/2007	Branscome et al.	N/A	N/A
2008/0243675	12/2007	Parsons et al.	N/A	N/A
2008/0253395	12/2007	Pandya	N/A	N/A
2008/0275805	12/2007	Hecht	N/A	N/A
2009/0037514	12/2008	Lankford et al.	N/A	N/A
2009/0182683	12/2008	Taylor et al.	N/A	N/A
2009/0262741	12/2008	Jungck et al.	N/A	N/A
2009/0287628	12/2008	Indeck et al.	N/A	N/A
2010/0005036	12/2009	Kraus et al.	N/A	N/A
2010/0027545	12/2009	Gomes et al.	N/A	N/A
2010/0082895	12/2009	Branscome et al.	N/A	N/A
2010/0106976	12/2009	Aciicmez et al.	N/A	N/A
2010/0174770	12/2009	Pandya	N/A	N/A
2010/0198920	12/2009	Wong et al.	N/A	N/A
2010/0257537	12/2009	Hinshaw et al.	N/A	N/A

2010/0306479	12/2009	Ezzat	N/A	N/A
2011/0029471	12/2010	Chakradhar et al.	N/A	N/A
2011/0040701	12/2010	Singla et al.	N/A	N/A
2011/0040776	12/2010	Najm et al.	N/A	N/A
2011/0066832	12/2010	Casselman et al.	N/A	N/A
2011/0125960	12/2010	Casselman	N/A	N/A
2011/0145130	12/2010	Glodjo et al.	N/A	N/A
2011/0167083	12/2010	Branscome et al.	N/A	N/A
2011/0178911	12/2010	Parsons et al.	N/A	N/A
2011/0178912	12/2010	Parsons et al.	N/A	N/A
2011/0178917	12/2010	Parsons et al.	N/A	N/A
2011/0178918	12/2010	Parsons et al.	N/A	N/A
2011/0178919	12/2010	Parsons et al.	N/A	N/A
2011/0178957	12/2010	Parsons et al.	N/A	N/A
2011/0179050	12/2010	Parsons et al.	N/A	N/A
2011/0184844	12/2010	Parsons et al.	N/A	N/A
2011/0199243	12/2010	Fallon et al.	N/A	N/A
2011/0218987	12/2010	Branscome et al.	N/A	N/A
2011/0231446	12/2010	Buhler et al.	N/A	N/A
2011/0246353	12/2010	Kraus et al.	N/A	N/A
2011/0252008	12/2010	Chamberlain et al.	N/A	N/A
2011/0289230	12/2010	Ueno	N/A	N/A
2011/0295967	12/2010	Wang et al.	N/A	N/A
2012/0065956	12/2011	Irturk et al.	N/A	N/A
2012/0089496	12/2011	Taylor et al.	N/A	N/A
2012/0089497	12/2011	Taylor et al.	N/A	N/A
2012/0095893	12/2011	Taylor et al.	N/A	N/A
2012/0109849	12/2011	Chamberlain et al.	N/A	N/A
2012/0110316	12/2011	Chamberlain et al.	N/A	N/A
2012/0116998	12/2011	Indeck et al.	N/A	N/A
2012/0117610	12/2011	Pandya	N/A	N/A
2012/0130922	12/2011	Indeck et al.	N/A	N/A
2012/0179590	12/2011	Borkovec et al.	N/A	N/A
2012/0215801	12/2011	Indeck et al.	N/A	N/A
2012/0246052	12/2011	Taylor et al.	N/A	N/A
2013/0007000	12/2012	Indeck et al.	N/A	N/A
2013/0018835	12/2012	Pandya	N/A	N/A
2013/0086096	12/2012	Indeck et al.	N/A	N/A
2013/0151458	12/2012	Indeck et al.	N/A	N/A
2013/0159449	12/2012	Taylor et al.	N/A	N/A
2013/0202287	12/2012	Joffe	N/A	N/A
2013/0262287	12/2012	Parsons et al.	N/A	N/A
2013/0290163	12/2012	Parsons et al.	N/A	N/A
2014/0025656	12/2013	Indeck et al.	N/A	N/A
2014/0040109	12/2013	Parsons et al.	N/A	N/A
2014/0067830	12/2013	Buhler et al.	N/A	N/A
2014/0089163	12/2013	Parsons et al.	N/A	N/A
2014/0164215	12/2013	Parsons et al.	N/A	N/A
2014/0180903	12/2013	Parsons et al.	N/A	N/A
2014/0180904	12/2013	Parsons et al.	N/A	N/A
2014/0180905	12/2013	Parsons et al.	N/A	N/A
2014/0181133	12/2013	Parsons et al.	N/A	N/A
2014/0310148	12/2013	Taylor et al.	N/A	N/A

2014/0310717	12/2013	Chamberlain et al.	N/A	N/A
2016/0070583	12/2015	Chamberlain et al.	N/A	N/A
2017/0124255	12/2016	Buhler et al.	N/A	N/A
2018/0276271	12/2017	Parsons et al.	N/A	N/A
2018/0330444	12/2017	Taylor et al.	N/A	N/A
2021/0142218	12/2020	Chamberlain et al.	N/A	N/A
2021/0200559	12/2020	Chamberlain et al.	N/A	N/A
2021/0304848	12/2020	Buhler et al.	N/A	N/A
2022/0358592	12/2021	Taylor et al.	N/A	N/A
2022/0414778	12/2021	Parsons et al.	N/A	N/A
2024/0037660	12/2023	Taylor et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
0573991	12/1992	EP	N/A
0880088	12/1995	EP	N/A
09-269901	12/1996	EP	N/A
0851358	12/1997	EP	N/A
0887723	12/1997	EP	N/A
0911738	12/1998	EP	N/A
09145544	12/1996	JP	N/A
11-259559	12/1998	JP	N/A
11282912	12/1998	JP	N/A
11316765	12/1998	JP	N/A
2000286715	12/1999	JP	N/A
2001268071	12/2000	JP	N/A
2001283000	12/2000	JP	N/A
2002101089	12/2001	JP	N/A
2002269343	12/2001	JP	N/A
2002352070	12/2001	JP	N/A
2003-036360	12/2002	JP	N/A
2003256660	12/2002	JP	N/A
2006059203	12/2005	JP	N/A
2006293852	12/2005	JP	N/A
4180644	12/2007	JP	N/A
199010910	12/1989	WO	N/A
199409443	12/1993	WO	N/A
199737735	12/1996	WO	N/A
2000041136	12/1999	WO	N/A
2001022425	12/2000	WO	N/A
0135216	12/2000	WO	N/A
200172106	12/2000	WO	N/A
2001080082	12/2000	WO	N/A
2001080558	12/2000	WO	N/A
0190890	12/2000	WO	N/A
2002061525	12/2001	WO	N/A
2003100650	12/2002	WO	N/A
2003036845	12/2002	WO	N/A
2003100662	12/2002	WO	N/A
2003104943	12/2002	WO	N/A
2004017604	12/2003	WO	N/A
2004042560	12/2003	WO	N/A
2004042561	12/2003	WO	N/A

2004042562	12/2003	WO	N/A
2004042574	12/2003	WO	N/A
2005017708	12/2004	WO	N/A
2005026925	12/2004	WO	N/A
2005048134	12/2004	WO	N/A
2005081855	12/2004	WO	N/A
2005114339	12/2004	WO	N/A
2006023948	12/2005	WO	N/A
2006060571	12/2005	WO	N/A
2006096324	12/2005	WO	N/A
2007064685	12/2006	WO	N/A
2007/074903	12/2006	WO	N/A
2007079095	12/2006	WO	N/A
2007087507	12/2006	WO	N/A
2007/127336	12/2006	WO	N/A
2008022036	12/2007	WO	N/A
2008073824	12/2007	WO	N/A
2009089467	12/2008	WO	N/A
2009140363	12/2008	WO	N/A
2010/077829	12/2009	WO	N/A
2013090363	12/2012	WO	N/A
2013/148693	12/2012	WO	N/A

OTHER PUBLICATIONS

Jones et al., “A Probabilistic Model of Information Retrieval: Development and Status”, Information Processing and Management, Aug. 1998, 76 pages. cited by applicant

Keutzer et al., “A Survey of Programmable Platforms—Network Proc”, University of California-Berkeley, pp. 1-29. cited by applicant

Koloniari et al., “Content-Based Routing of Path Queries in Peer-to-Peer Systems”, pp. 1-19, E. Bertino et al. (Eds.): EDBT 2004, LNCS 2992, pp. 29-47, 2004, copyright by Springer-Verlag, Germany. cited by applicant

Krishnamurthy et al., “Biosequence Similarity Search On The Mercury System”, Proceedings of the 15th IEEE International Conference on Application-Specific Systems, Architectures, and Processors (ASAP04), Sep. 2004, pp. 365-375. cited by applicant

Kulig et al., “System and Method for Controlling Transmission of Data Packets Over an Information Network”, pending U.S. Patent Application. cited by applicant

Lancaster et al., “Acceleration of Ungapped Extension in Mercury BLAST”, Seventh (7th) Workshop on Media and Streaming Processors, Nov. 12, 2005, Thirty-Eighth (38th) International Symposium on Microarchitecture (MICRO-38), Barcelona, Spain. cited by applicant

Li et al., “Large-Scale IP Traceback in High-Speed Internet: Practical Techniques and Theoretical Foundation”, Proceedings of the 2004 IEEE Symposium on Security and Privacy, 2004, pp. 1-15. cited by applicant

Lin et al., “Real-Time Image Template Matching Based on Systolic Array Processor”, International Journal of Electronics; Dec. 1, 1992; pp. 1165-1176; vol. 73, No. 6; London, Great Britain. cited by applicant

Lockwood et al., “Field Programmable Port Extender (FPX) for Distributed Routing and Queuing”, ACM International Symposium on Field Programmable Gate Arrays (FPGA 2000), Monterey, CA, Feb. 2000, pp. 137-144. cited by applicant

Lockwood et al., “Hello, World: A Simple Application for the Field Programmable Port Extender (FPX)”, Washington University, Department of Computer Science, Technical Report WUCS-00-12, Jul. 11, 2000. cited by applicant

Lockwood et al., “Parallel FPGA Programming over Backplane Chassis”, Washington University, Department of Computer Science, Technical Report WUCS-00-11, Jun. 12, 2000. cited by applicant

Lockwood et al., “Reprogrammable Network Packet Processing on the Field Programmable Port Extender (Fpx)”, ACM International Symposium on Field Programmable Gate Arrays (FPGA 2001), Monterey, CA, Feb. 2001, pp. 87-93. cited by applicant

Lockwood, "An Open Platform for Development of Network Processing Modules in Reprogrammable Hardware", IEC DesignCon 2001, Santa Clara, CA, Jan. 2001, Paper WB-19. cited by applicant

Lockwood, "Building Networks with Reprogrammable Hardware", Field Programmable Port Extender: Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002. cited by applicant

Lockwood, "Evolvable Internet Hardware Platforms", NASA/DoD Workshop on Evolvable Hardware (EHW'01), Long Beach, CA, Jul. 12-14, 2001, pp. 271-279. cited by applicant

Lockwood, "Hardware Laboratory Configuration", Field Programmable Port Extender: Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002. cited by applicant

Lockwood, "Introduction", Field Programmable Port Extender: Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002. cited by applicant

Lockwood, "Platform and Methodology for Teaching Design of Hardware Modules in Internet Routers and Firewalls", IEEE Computer Society International Conference on Microelectronic Systems Education (MSE'2001), Las Vegas, NV, Jun. 17-18, 2001, pp. 56-57. cited by applicant

Lockwood, "Protocol Processing on the FPX", Field Programmable Port Extender: Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002. cited by applicant

Lockwood, "Simulation and Synthesis", Field Programmable Port Extender: Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002. cited by applicant

Lockwood, "Simulation of the Hello World Application for the Field-Programmable Port Extender (FPX)", Washington University, Applied Research Lab, Spring 2001 Gigabits Kits Workshop. cited by applicant

Madhusudan, "Design of a System for Real-Time Worm Detection", Hot Interconnects, pp. 77-83, Stanford, CA, Aug. 2004, found at <http://www.hoti.org/hoti12/program/papers/2004/paper4.2.pdf>. cited by applicant

Madhusudan, "Design of a System for Real-Time Worm Detection", Power Point Presentation in Support of Master's Thesis, Washington Univ., Dept. of Computer Science and Engineering, St. Louis, MO, Aug. 2004. cited by applicant

Mao et al., "Cluster-based Online Monitoring System of Web Traffic", Dept. of Computer Science and Technology, Tsinghua Univ., Beijing, 100084 P.R. China. cited by applicant

Mosanya et al., "A FPGA-Based Hardware Implementation of Generalized Profile Search Using Online Arithmetic", ACM/Sigda International Symposium on Field Programmable Gate Arrays (FPGA '99), Feb. 21-23, 1999, pp. 101-111, Monterey, CA, USA. cited by applicant

Moscola et al., "FPGrep and FPSed: Regular Expression Search and Substitution for Packet Streaming in Field Programmable Hardware", Dept. of Computer Science, Applied Research Lab, Washington University, Jan. 8, 2002, unpublished, pp. 1-19, St. Louis, MO. cited by applicant

Moscola et al., "FPSed: A Streaming Content Search-and-Replace Module for an Internet Firewall", Proc. of Hot Interconnects, 11th Symposium on High Performance Interconnects, pp. 122-129, Aug. 20, 2003. cited by applicant

Moscola, "FPGrep and FPSed: Packet Payload Processors for Managing the Flow of Digital Content on Local Area Networks and the Internet", Master's Thesis, Sever Institute of Technology, Washington University, St. Louis, MO, Aug. 2003. cited by applicant

Motwani et al., "Randomized Algorithms", 1995, pp. 215-216, Cambridge University Press. cited by applicant

Mueller, "Upgrading and Repairing PCs, 15th Anniversary Edition", 2004, pp. 63-66, 188, Que. cited by applicant

Navarro, "A Guided Tour to Approximate String Matching", ACM Computing Surveys, vol. 33, No. 1, Mar. 2001, pp. 31-88. cited by applicant

Necker et al., "TCP-Stream Reassembly and State Tracking in Hardware", School of Electrical and Computer Engineering, Georgia Institute of Technology, Atlanta, GA. cited by applicant

Nunez et al., "The X-MatchLITE FPGA-Based Data Compressor", Euromicro Conference 1999, Proceedings, Italy, Sep. 8-10, 1999, pp. 126-132, Los Alamitos, CA. cited by applicant

Nwodoh et al., "A Processing System for Real-Time Holographic Video Computation", Reconfigurable Technology: FPGAs for Computing and Application, Proceedings for the SPIE, Sep. 1999, Boston, pp. 129-140, vol. 3844. cited by applicant

Office Action for EP Application 13767579.9 dated Sep. 6, 2016. cited by applicant

Office Action for U.S. Appl. No. 14/195,550 dated Jun. 27, 2016. cited by applicant

Office Action for U.S. Appl. No. 14/195,462 dated Aug. 10, 2017. cited by applicant

Office Action for U.S. Appl. No. 14/195,462 dated Aug. 26, 2016. cited by applicant
Office Action for U.S. Appl. No. 14/195,462 dated Feb. 12, 2016. cited by applicant
Office Action for U.S. Appl. No. 14/195,462 dated Jan. 27, 2017. cited by applicant
Office Action for U.S. Appl. No. 14/195,510 dated Mar. 10, 2017. cited by applicant
Office Action for U.S. Appl. No. 14/195,510 dated Nov. 30, 2017. cited by applicant
Office Action for U.S. Appl. No. 14/195,531 dated Feb. 24, 2017. cited by applicant
Office Action for U.S. Appl. No. 14/195,531 dated May 18, 2016. cited by applicant
Office Action for U.S. Appl. No. 14/195,531 dated Sep. 7, 2017. cited by applicant
Office Action for U.S. Appl. No. 14/195,550 dated Feb. 10, 2017. cited by applicant
Parsing—(Wikipedia definition) <https://en.wikipedia.org/w/index.php?title=Parsing&oldid=475958964>. cited by applicant
Prakash et al., “OC-3072 Packet Classification Using BDDs and Pipelined SRAMs”, Department of Electrical and Computer Engineering, The University of Texas at Austin. cited by applicant
Pramanik et al., “A Hardware Pattern Matching Algorithm on a Dataflow”; Computer Journal; Jul. 1, 1985; pp. 264-269; vol. 28, No. 3; Oxford University Press, Surrey, Great Britain. cited by applicant
“A Reconfigurable Computing Model for Biological Research Application of Smith-Waterman Analysis to Bacterial Genomes” A White Paper Prepared by Star Bridge Systems, Inc. [retrieved Dec. 12, 2006]. Retrieved from the Internet: <URL: <http://www.starbridgesystems.com/resources/whitepapers/Smith%20Waterman%20Whitepaper.pdf>. cited by applicant
“ACTIV Financial Announces Hardware Based Market Data Feed Processing Strategy”, For Release on Apr. 2, 2007, 2 pages. cited by applicant
“ACTIV Financial Delivers Accelerated Market Data Feed”, Apr. 6, 2007, byline of Apr. 2, 2007, downloaded from <http://hpcwire.com/hpc.1346816.html> on Jun. 19, 2007, 3 pages. cited by applicant
“Curriculum Vitae of Michael Brogioli, Ph.D” for Inter Partes Review Petition re U.S. Pat. No. 10,121,196; pp. 1-19; dated May 28, 2020. cited by applicant
“Declaration of Michael Brogioli, Ph.D” for Inter Partes Review Petition re U.S. Pat. No. 10,121,196; pp. 1-112; dated May 28, 2020. cited by applicant
“DRC, Exegy Announce Joint Development Agreement”, Jun. 8, 2007, byline of Jun. 4, 2007; downloaded from <http://www.hpcwire.com/hpc/1595644.html> on Jun. 19, 2007, 3 pages. cited by applicant
“Lucent Technologies Delivers “PayloadPlus” Network Processors for Programmable, MultiProtocol, OC-48c Processing”, Lucent Technologies Press Release, downloaded from <http://www.lucent.com/press/1000/0010320.meb.html> on Mar. 21, 2002. cited by applicant
“Overview, Field Programmable Port Extender”, Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002, pp. 1-4. cited by applicant
“Payload Plus™ Agere System Interface”, Agere Systems Product Brief, Jun. 2001, downloaded from Internet, Jan. 2002, pp. 1-6. cited by applicant
“Petition for Inter Partes Review” re U.S. Pat. No. 10,121,196; pp. 1-89; dated May 28, 2020. cited by applicant
“RFC793: Transmission Control Protocol, Darpa Internet Program, Protocol Specification”, Sep. 1981. cited by applicant
“Technology Overview”, Data Search Systems Incorporated, downloaded from the <http://www.datasearchsystems.com/tech.htm> on Apr. 19, 2004. cited by applicant
“The Field-Programmable Port Extender (FPX)”, downloaded from <http://www.arl.wustl.edu/arl/> in Mar. 2002. cited by applicant
Aldwairi et al., “Configurable String Matching Hardware for Speeding up Intrusion Detection”, SIRARCH Comput. Archit. News, vol. 33, No. 1, pp. 99-107, Mar. 2005. cited by applicant
Amanuma et al., “A FPGA Architecture For High Speed Computation”, Proceedings of 60th Convention Architecture, Software Science, Engineering, Mar. 14, 2000, pp. 1-163-1-164, Information Processing Society, Japan. cited by applicant
Anerousis et al., “Using the AT&T Labs PacketScope for Internet Measurement, Design, and Performance Analysis”, Network and Distributed Systems Research Laboratory, AT&T Labs-Research, Florham, Park, NJ, Oct. 1997. cited by applicant
Anonymous, “Method for Allocating Computer Disk Space to a File of Known Size”, IBM Technical

Disclosure Bulletin, vol. 27, No. 10B, Mar. 1, 1985, New York. cited by applicant

Apostolopoulos et al., "Design, Implementaion and Performance of a Content-Based Switch", Proceedings IEEE INFOCOM 2000, The Conference on Computer Communications, 19th Annual Joint Conference of the IEEE Computer and Communications Societes, Tel Aviv, Israel, Mar. 26-30, 2000, pp. 1117-1126. cited by applicant

Arnold et al., "The Splash 2 Processor and Applications", Proceedings 1993 IEEE International Conference on Computer Design: VLSI In Computers and Processors (ICCD '93), Oct. 3, 1993, pp. 482-485, IEEE Computer Society, Cambridge, MA USA. cited by applicant

Artan et al., "Multi-packet Signature Detection using Prefix Bloom Filters", 2005, IEEE, pp. 1811-1816. cited by applicant

Asami et al., "Improvement of DES Key Search on FPGA-Based Parallel Machine "RASH"", Proceedings of Information Processing Society, Aug. 15, 2000, pp. 50-57, vol. 41, No. SIG5 (HPS1), Japan. cited by applicant

Baboescu et al., "Scalable Packet Classification," SIGCOMM'01, Aug. 27-31, 2001, pp. 199-210, San Diego, California, USA; <http://www.ecse.rpi.edu/homepages/shivkuma/teaching/sp2001/readings/baboescu-pkt-classification.pdf>. cited by applicant

Baer, "Computer Systems Architecture", 1980, pp. 262-265; Computer Science Press, Potomac, Maryland. cited by applicant

Baeza-Yates et al., "New and Faster Filters for Multiple Approximate String Matching", Random Structures and Algorithms (RSA), Jan. 2002, pp. 23-49, vol. 20, No. 1. cited by applicant

Baker et al., "High-throughput Linked-Pattern Matching for Intrusion Detection Systems", ANCS 2005: Proceedings of the 2005 Symposium on Architecture for Networking and Communications Systems, pp. 193-202, ACM Press, 2005. cited by applicant

Barone-Adesi et al., "Efficient Analytic Approximation of American Option Values", Journal of Finance, vol. 42, No. 2 (Jun. 1987), pp. 301-320. cited by applicant

Batory, "Modeling the Storage Architectures of Commercial Database Systems", ACM Transactions on Database Systems, Dec. 1985, pp. 463-528, vol. 10, issue 4. cited by applicant

Behrens et al., "BLASTN Redundancy Filter in Reprogrammable Hardware," Final Project Submission, Fall 2003, Department of Computer Science and Engineering, Washington University. cited by applicant

Berk, "JLex: A lexical analyzer generator for Java™", downloaded from <http://www.cs.princeton.edu/~appel/modern/java/Jlex/> in Jan. 2002, pp. 1-18. cited by applicant

Bianchi et al., "Improved Queueing Analysis of Shared Buffer Switching Networks", ACM, Aug. 1993, pp. 482-490. cited by applicant

Bloom, "Space/Time Trade-offs in Hash Coding With Allowable Errors", Communications of the ACM, Jul. 1970, pp. 422-426, vol. 13, No. 7, Computer Usage Company, Newton Upper Falls, Massachusetts, USA. cited by applicant

Braun et al., "Layered Protocol Wrappers for Internet Packet Processing in Reconfigurable Hardware", Proceedings of Hot Interconnects 9 (HotI-9) Stanford, CA, Aug. 22-24, 2001, pp. 93-98. cited by applicant

Braun et al., "Protocol Wrappers for Layered Network Packet Processing in Reconfigurable Hardware", IEEE Micro, Jan.-Feb. 2002, pp. 66-74. cited by applicant

Brodie et al., "Dynamic Reconfigurable Computing", in Proc. of 9th Military and Aerospace Programmable Logic Devices International Conference, Sep. 2006. cited by applicant

Cavnar et al., "N-Gram-Based Text Categorization", Proceedings of SDAIR-94, 3rd Annual Symposium on Document Analysis and Information Retrieval, Las Vegas, pp. 161-175, 1994. cited by applicant

Celko, "Joe Celko's Data & Databases: Concepts in Practice", 1999, pp. 72-74, Morgan Kaufmann Publishers. cited by applicant

Chamberlain et al., "Achieving Real Data Throughput for an FPGA Co-Processor on Commodity Server Platforms", Proc. of 1st Workshop on Building Block Engine Architectures for Computers and Networks, Oct. 2004, Boston, MA. cited by applicant

Chamberlain et al., "The Mercury System: Embedding Computation Into Disk Drives", 7th High Performance Embedded Computing Workshop, Sep. 2003, Boston, MA. cited by applicant

Chamberlain et al., "The Mercury System: Exploiting Truly Fast Hardware for Data Search", Proc. of Workshop on Storage Network Architecture and Parallel I/Os, Sep. 2003, New Orleans, LA. cited by applicant

Chaney et al., "Design of a Gigabit ATM Switch", Washington University, St. Louis. cited by applicant

Cho et al., “Deep Packet Filter with Dedicated Logic and Read Only Memories”, 12th Annual IEEE Symposium on Field-Programmable Custom Computing Machines, Apr. 2004. cited by applicant

Choi et al., “Design of a Flexible Open Platform for High Performance Active Networks”, Allerton Conference, 1999, Champaign, IL. cited by applicant

Cholleti, “Storage Allocation in Bounded Time”, MS Thesis, Dept. of Computer Science and Engineering, Washington University, St. Louis, MO (Dec. 2002). Available as Washington University Technical Report WUCSE-2003-2. cited by applicant

Clark et al., “Scalable Pattern Matching for High Speed Networks”, Proceedings of the 12th Annual IEEE Symposium on Field-Programmable Custom Computing Machines, 2004; FCCM 2004, Apr. 20-23, 2004; pp. 249-257; IEEE Computer Society; Cambridge, MA USA. cited by applicant

Cloutier et al., “VIP: An FPGA-Based Processor for Image Processing and Neural Networks”, Proceedings of Fifth International Conference on Microelectronics for Neural Networks, Feb. 12, 1996, pp. 330-336, Los Alamitos, California. cited by applicant

Compton et al., “Configurable Computing: A Survey of Systems and Software”, Technical Report, Northwestern University, Dept. of ECE, 1999. cited by applicant

Compton et al., “Reconfigurable Computing: A Survey of Systems and Software”, Technical Report, Northwestern University, Dept. of ECE, 1999, presented by Yi-Gang Tai. cited by applicant

Compton et al., “Reconfigurable Computing: A Survey of Systems and Software”, University of Washington, ACM Computing Surveys, Jun. 2, 2002, pp. 171-210, vol. 34 No. 2, <<http://www.idi.ntnu.no/emner/tdt22/2011/reconfig.pdf>>. cited by applicant

Cong et al., “An Optional Technology Mapping Algorithm for Delay Optimization in Lookup-Table Based FPGA Designs”, IEEE, 1992, pp. 48-53. cited by applicant

Crosman, “Who Will Cure Your Data Latency?”, Storage & Servers, Jun. 20, 2007, URL: <http://www.networkcomputing.com/article/printFullArticleSrc.jhtml?articleID=199905630>. cited by applicant

Cuppu and Jacob, “Organizational Design Trade-Offs at the DRAM, Memory Bus and Memory Controller Level: Initial Results,” Technical Report UMB-SCA-1999-2, Univ. of Maryland Systems & Computer Architecture Group, Nov. 1999, pp. 1-10. cited by applicant

Currid, “TCP Offload to the Rescue”, Networks, Jun. 14, 2004, 16 pages, vol. 2, No. 3. cited by applicant

Data Structure—(Wikipedia definition) https://en.wikipedia.org/wiki/Data_structure (4 pages). cited by applicant

Denoyer et al., “HMM-based Passage Models for Document Classification and Ranking”, Proceedings of ECIR-01, 23rd European Colloquium Information Retrieval Research, Darmstadt, DE, pp. 126-135, 2001. cited by applicant

Dharmapurikar et al., “Deep Packet Inspection Using Parallel Bloom Filters,” IEEE Micro, Jan.-Feb. 2004, vol. 24, Issue: 1, pp. 52-61. cited by applicant

Dharmapurikar et al., “Deep Packet Inspection Using Parallel Bloom Filters,” Symposium on High Performance Interconnects (HotI), Stanford, California, 2003, pp. 44-51. cited by applicant

Dharmapurikar et al., “Design and Implementation of a String Matching System for Network Intrusion Detection using FPGA-based Bloom Filters”, Proc. of 12th Annual IEEE Symposium on Field Programmable Custom Computing Machines, 2004, pp. 1-10. cited by applicant

Dharmapurikar et al., “Longest Prefix Matching Using Bloom Filters,” SIGCOMM, 2003, pp. 201-212. cited by applicant

Dharmapurikar et al., “Robust TCP Stream Reassembly in the Presence of Adversaries”, Proc. of the 14th Conference on USENIX Security Symposium—vol. 14, 16 pages, Baltimore, MD, 2005; <http://www.icir.org/vern/papers/TcpReassembly/TCPReassembly.pdf>. cited by applicant

Dharmapurikar, “Fast and Scalable Pattern Matching for Content Filtering”, ACM, ANCS 05, 2005, pp. 183-192. cited by applicant

Ebeling et al., “RaPiD—Reconfigurable Pipelined Datapath”, University of Washington, Dept. of Computer Science and Engineering, Sep. 23, 1996, Seattle, WA. cited by applicant

Exegy Inc., “Exegy and HyperFeed to Unveil Exelerate TP at SIA Conference”, Release Date: Jun. 20, 2006, downloaded from <http://news.thomasnet.com/companystory/488004> on Jun. 19, 2007, 4 pages. cited by applicant

Exegy Inc., “First Exegy Ticker Plant Deployed”, Release Date: Oct. 17, 2006, downloaded from

<http://news.companystory/496530> on Jun. 19, 2007, 5 pages. cited by applicant

Extended European Search Report for EP Application 11847815.5 dated Apr. 4, 2014. cited by applicant

Extended European Search Report for EP Application 13767579.9 dated Oct. 22, 2015. cited by applicant

Feldman, "High Frequency Traders Get Boost From FPGA Acceleration", Jun. 8, 2007, downloaded from <http://www.hpcwire.com/hpc.1600113.html> on Jun. 19, 2007, 4 pages. cited by applicant

Feldmann, "BLT: Bi-Layer Tracing of HTTP and TCP/IP", AT&T Labs-Research, Florham Park, NJ, USA. cited by applicant

Franklin et al., "An Architecture for Fast Processing of Large Unstructured Data Sets." Proc. of 22nd Int'l Conf. on Computer Design, Oct. 2004, pp. 280-287. cited by applicant

Franklin et al., "Assisting Network Intrusion Detection with Reconfigurable Hardware", Symposium on Field-Programmable Custom Computing Machines (FCCM 2002), Apr. 2002, Napa, California. cited by applicant

Fu et al., "The Fpx Kcpsm Module: An Embedded, Reconfigurable Active Processing Module for the Field Programmable Port Extender (FPX)", Washington University, Department of Computer Science, Technical Report WUCS-01-14, Jul. 2001. cited by applicant

Gavrila et al., "Multi-feature Hierarchical Template Matching Using Distance Transforms", IEEE, Aug. 16-20, 1998, vol. 1, pp. 439-444. cited by applicant

Gokhale et al., "Reconfigurable Computing: Accelerating Computation With Field-Programmable Gate Arrays", 2005, pp. 1-3, 7, 11-15, 39, 92-93, Springer. cited by applicant

Gokhale et al., "Reconfigurable Computing: Accelerating Computation with Field-Programmable Gate Arrays", Springer, 2005, pp. 1-36. cited by applicant

Gokhale et al., "Reconfigurable Computing: Accelerating Computation with Field-Programmable Gate Arrays", Springer, 2005, pp. 1-54, 92-96. cited by applicant

Google Search Results Page for "field programmable gate array financial calculation stock market" over dates of Jan. 1, 1990-May 21, 2002, 1 page. cited by applicant

Gunther et al., "Assessing Document Relevance with Run-Time Reconfigurable Machines", IEEE Symposium on FPGAs for Custom Computing Machines, 1996, pp. 10-17, Proceedings, Napa Valley, CA. cited by applicant

Gupta et al., "High-Speed Implementations of Rule-Based Systems," ACM Transactions on Computer Systems, May 1989, pp. 119-146, vol. 7, Issue 2. cited by applicant

Gupta et al., "Packet Classification on Multiple Fields", Computer Systems Laboratory, Stanford University, Stanford, CA. cited by applicant

Gupta et al., "PMM: A Parallel Architecture for Production Systems," Proceedings of the IEEE, Apr. 1992, pp. 693-696, vol. 2. cited by applicant

Gurtov, "Effect of Delays on TCP Performance", Cellular Systems Development, Sonera Corporation, online at <http://cs.helsinki.fi/u/gurtov/papers/pwc01.pdf>. cited by applicant

Gyang, "NCBI BLASTN Stage 1 in Reconfigurable Hardware," Technical Report WUCSE-2005-30, Aug. 2004, Department of Computer Science and Engineering, Washington University, St. Louis, MO. cited by applicant

Halaas et al., "A Recursive MISD Architecture for Pattern Matching", IEEE Transactions on Very Large Scale Integration, vol. 12, No. 7, pp. 727-734, Jul. 2004. cited by applicant

Harris, "Pete's Blog: Can FPGAs Overcome the FUD?", Low-Latency.com, May 14, 2007, URL: <http://www.a-teamgroup.com/article/pete-blog-can-fpgas-overcome-the-fud/>. cited by applicant

Hauck et al., "Software Technologies for Reconfigurable Systems", Northwestern University, Dept. of ECE, Technical Report, 1996. cited by applicant

Hayes, "Computer Architecture and Organization", Second Edition, 1988, pp. 448-459, McGraw-Hill, Inc. cited by applicant

Hezel et al., "FPGA-Based Template Matching Using Distance Transforms", Proceedings of the 10th Annual IEEE Symposium on Field-Programmable Custom Computing Machines, Apr. 22, 2002, pp. 89-97, IEEE Computer Society, USA. cited by applicant

Hirsch, "Tech Predictions for 2008", Reconfigurable Computing, Jan. 16, 2008, URL: http://fpgacomputing.blogspot.com/2008_01_01_archive.html. cited by applicant

Hoinville, et al. "Spatial Noise Phenomena of Longitudinal Magnetic Recording Media", IEEE Transactions on Magnetics, vol. 28, No. 6, Nov. 1992. cited by applicant

Hollaar, "Hardware Systems for Text Information Retrieval", Proceedings of the Sixth Annual International ACM Sigir Conference on Research and Development in Information Retrieval, Jun. 6-8, 1983, pp. 3-9, Baltimore, Maryland, USA. cited by applicant

Howe, Data Analysis for Database Design Third Edition, 2001, 335 pages, Butterworth-Heinemann. cited by applicant

Hutchings et al., "Assisting Network Intrusion Detection with Reconfigurable Hardware", FCCM 2002: 10th Annual IEEE Symposium on Field-Programmable Custom Computing Machines, 2002. cited by applicant

Ibrahim et al., "Lecture Notes in Computer Science: Database and Expert Systems Applications", 2000, p. 769, vol. 1873, Springer. cited by applicant

International Preliminary Report on Patentability (Chapter I) for PCT/US2011/064269 issued Jun. 12, 2013. cited by applicant

International Preliminary Report on Patentability (Chapter I) for PCT/US2013/033889 issued Oct. 9, 2014. cited by applicant

International Search Report and Written Opinion for PCT/US2011/064269 dated Apr. 20, 2012. cited by applicant

International Search Report and Written Opinion for PCT/US2012/069142 dated Feb. 22, 2013. cited by applicant

International Search Report and Written Opinion for PCT/US2013/033889 dated Aug. 29, 2013. cited by applicant

Jacobson et al., "RFC 1072: TCP Extensions for Long-Delay Paths", Oct. 1988. cited by applicant

Jacobson et al., "tcpdump—dump traffic on a network", Jun. 30, 1997, online at www.cse.cuhk.edu.hk/~cslui/CEG4430/tcpdump.ps.gz. cited by applicant

Johnson et al., "Pattern Matching in Reconfigurable Logic for Packet Classification", College of Computing, Georgia Institute of Technology, Atlanta, GA. cited by applicant

"Declaration of David E. Taylor" for Inter Partes Review re U.S. Pat. No. 10,121,196; IPR2020-01006, pp. 1-11; dated Sep. 8, 2020. cited by applicant

"Declaration of Robert Zeidman" for Inter Partes Review re U.S. Pat. No. 10,121,196; IPR2020-01006, pp. 1-74; dated Feb. 28, 2021. cited by applicant

"Deposition Transcript of Michael C. Brogioli, Ph.D." for Inter Partes Review re U.S. Pat. No. 10,121,196; IPR2020-01006, pp. 1-74; dated Feb. 9, 2021. cited by applicant

"Patent Owner's Preliminary Response" for Inter Partes Review re U.S. Pat. No. 10,121,196; IPR2020-01006, pp. 1-28; dated Sep. 8, 2020. cited by applicant

"Patent Owner's Response" for Inter Partes Review re U.S. Pat. No. 10,121,196; IPR2020-01006, pp. 1-74; dated Mar. 1, 2021. cited by applicant

"Second Declaration of David E. Taylor" for Inter Partes Review re U.S. Pat. No. 10,121,196; IPR2020-01006, pp. 1-15; dated Feb. 27, 2021. cited by applicant

Diniz et al., "Data Search and Reorganization Using FPGAs: Application to Spatial Pointer-Based Data Structures", IEEE, 2003, 11 pgs. cited by applicant

Ferguson, "Vector Signal Analysis of Digital Baseband and if Signals within an FPGA", 2005, IEEE, pp. 402-407. cited by applicant

Gokhale et al., "Reconfigurable Computing: Accelerating Computation with Field-Programmable Gate Arrays", Springer, 2005, pp. 1-209. cited by applicant

Prosecution History of U.S. Appl. No. 14/195,531, now U.S. Pat. No. 11,436,672, filed Mar. 3, 2014. cited by applicant

Villasenor et al., "The Flexibility of Configurable Computing", IEEE, 1998, pp. 67-84. cited by applicant

Vuillemin et al., "Programmable Active Memories: Reconfigurable Systems Come of Age", IEEE, 1996, pp. 56-69, vol. 4, No. 1. cited by applicant

Prosecution History for U.S. Appl. No. 14/195,550, now U.S. Pat. No. 9,990,393, filed Mar. 3, 2014. cited by applicant

Prosecution History for U.S. Appl. No. 13/833,098, filed Mar. 15, 2013, now U.S. Pat. No. 10,121,196, issued Nov. 6, 2018. cited by applicant

Ramakrishna et al., "A Performance Study of Hashing Functions for Hardware Applications", Int. Conf. on Computing and Information, May 1994, pp. 1621-1636, vol. 1, No. 1. cited by applicant

Ramakrishna et al., "Efficient Hardware Hashing Functions for High Performance Computers", IEEE Transactions on Computers, Dec. 1997, vol. 46, No. 12. cited by applicant

Ratha et al., "Convolution on Splash 2", Proceedings of IEEE Symposium on FPGAs for Custom Computing Machines, Apr. 19, 1995, pp. 204-213, Los Alamitos, California. cited by applicant

Roberts, "Internet Still Growing Dramatically Says Internet Founder", Press Release, Caspian Networks, Inc. —Virtual Pressroom. cited by applicant

Roesch, "Snort—Lightweight Intrusion Detection for Networks", Proceedings of LISA '99: 13th Systems Administration Conference; Nov. 7-12, 1999; pp. 229-238; USENIX Association, Seattle, WA USA. cited by applicant

Roy, "A bounded search algorithm for segmented channel routing for FPGA's and associated channel architecture issues", IEEE, Nov. 11, 1993, pp. 1695-1705, vol. 12. cited by applicant

Russ et al., Non-Intrusive Built-In Self-Test for FPGA and MCM Applications, Aug. 8-10, 1995, IEEE, 480-485. cited by applicant

Sachin Tandon, "A Programmable Architecture for Real-Time Derivative Trading", Master's Thesis, University of Edinburgh, 2003. cited by applicant

Schmerken, "With Hyperfeed Litigation Pending, Exegy Launches Low-Latency Ticker Plant", in Wall Street & Technology Blog, Mar. 20, 2007, pp. 1-2. cited by applicant

Schmit, "Incremental Reconfiguration for Pipelined Applications", FPGAs for Custom Computing Machines, Proceedings, The 5th Annual IEEE Symposium, Dept. of ECE, Carnegie Mellon University, Apr. 16-18, 1997, pp. 47-55, Pittsburgh, PA. cited by applicant

Schuehler et al., "Architecture for a Hardware Based, TCP/IP Content Scanning System", IEEE Micro, 24(1):62-69, Jan.-Feb. 2004, USA. cited by applicant

Schuehler et al., "TCP-Splitter: A TCP/IP Flow Monitor in Reconfigurable Hardware", Hot Interconnects 10 (HotI-10), Stanford, CA, Aug. 21-23, 2002, pp. 127-131. cited by applicant

Seki et al., "High Speed Computation of Shogi With FPGA", Proceedings of 58th Convention Architecture, Software Science, Engineering, Mar. 9, 1999, pp. 1-133-1-134. cited by applicant

Shah, "Understanding Network Processors", Version 1.0, University of California-Berkeley, Sep. 4, 2001. cited by applicant

Shalunov et al., "Bulk TCP Use and Performance on Internet 2", ACM SIGCOMM Internet Measurement Workshop, 2001. cited by applicant

Shasha et al., "Database Tuning", 2003, pp. 280-284, Morgan Kaufmann Publishers. cited by applicant

Shirazi et al., "Quantitative Analysis of FPGA-based Database Searching", Journal of VLSI Signal Processing Systems For Signal, Image, and Video Technology, May 2001, pp. 85-96, vol. 28, No. 1/2, Kluwer Academic Publishers, Dordrecht, NL. cited by applicant

Sidhu et al., "Fast Regular Expression Matching Using FPGAs", IEEE Symposium on Field Programmable Custom Computing Machines (FCCM 2001), Apr. 2001. cited by applicant

Sidhu et al., "String Matching on Multicontext FPGAs Using Self-Reconfiguration", FPGA '99: Proceedings of the 1999 ACM/SIGDA 7th International Symposium on Field Programmable Gate Arrays, Feb. 1999, pp. 217-226. cited by applicant

Singh et al., "The EarlyBird System for Real-Time Detection on Unknown Worms", Technical report CS2003-0761, Aug. 2003. cited by applicant

Skiena et al., "Programming Challenges: The Programming Contest Training Manual", 2003, pp. 30-31, Springer. cited by applicant

Sourdis and Pnevmatikatos, "Fast, Large-Scale String Match for a 10Gbps FPGA-based Network Intrusion Detection System", 13th International Conference on Field Programmable Logic and Applications, 2003. cited by applicant

Steinbach et al., "A Comparison of Document Clustering Techniques", KDD Workshop on Text Mining, 2000. cited by applicant

Tan et al., "A High Throughput String Matching Architecture for Intrusion Detection and Prevention", ISCA 2005: 32nd Annual International Symposium on Computer Architecture, pp. 112-122, 2005. cited by applicant

Taylor et al., "Dynamic Hardware Plugins (DHP): Exploiting Reconfigurable Hardware for High-Performance Programmable Routers", Computer Networks, 38(3): 295-310 (16), Feb. 21, 2002, and online at http://www.cc.gatech.edu/classes/AY2007/cs8803hpc_fall/papers/phplugins.pdf. cited by applicant

Taylor et al., “Generalized RAD Module Interface Specification of the Field Programmable Port Extender (FPX) Version 2”, Washington University, Department of Computer Science, Technical Report, Jul. 5, 2001, pp. 1-10. cited by applicant

Taylor et al., “Modular Design Techniques for the FPX”, Field Programmable Port Extender: Jan. 2002 Gigabit Workshop Tutorial, Washington University, St. Louis, MO, Jan. 3-4, 2002. cited by applicant

Taylor et al., “Scalable Packet Classification using Distributed Crossproducting of Field Labels”, Proceedings of IEEE Infocom, Mar. 2005, pp. 1-12, vol. 20, No. 1. cited by applicant

Taylor, “Models, Algorithms, and Architectures for Scalable Packet Classification”, doctoral thesis, Department of Computer Science and Engineering, Washington University, St. Louis, MO, Aug. 2004, pp. 1-201. cited by applicant

Thomson Reuters, “Mellanox InfiniBand Accelerates the Exegy Ticker Plant at Major Exchanges”, Jul. 22, 2008, URL: <http://www.reuters.com/article/pressRelease/idUS125385+22-Jul-2008+BW20080722>. cited by applicant

Uluski et al., “Characterizing Antivirus Workload Execution”, SIGARCH Comput. Archit. News, vol. 33, No. 1, pp. 90-98, Mar. 2005. cited by applicant

Villasenor et al., “Configurable Computing Solutions For Automatic Target Recognition”, FPGAs for Custom Computing Machines, 1996, Proceedings, IEEE Symposium on Napa Valley, CA, Apr. 17-19, 1996, pp. 70-79, 1996 IEEE, Napa Valley, CA, Los Alamitos, CA, USA. cited by applicant

Waldvogel et al., “Scalable High-Speed Prefix Matching”, ACM Transactions on Computer Systems, Nov. 2001, pp. 440-482, vol. 19, No. 4. cited by applicant

Ward et al., “Dynamically Reconfigurable Computing: A Novel Computation Technology with Potential to Improve National Security Capabilities”, May 15, 2003, A White Paper Prepared by Star Bridge Systems, Inc. [retrieved Dec. 12, 2006]. Retrieved from the Internet: <URL: <http://www.starbridgesystems.com/resources/whitepapers/Dynamically%20Reconfigurable%20Computing.pdf>>. cited by applicant

Weaver et al., “Very Fast Containment of Scanning Worms”, Proc. USENIX Security Symposium 2004, San Diego, CA, Aug. 2004, located at <http://www.icsi.berkeley.edu/~nweaver/containment/containment.pdf>. cited by applicant

Web-Pop (Professional Options Package) (www.pmpublishing.com). cited by applicant

West et al., “An FPGA-Based Search Engine for Unstructured Database”, Proc. of 2nd Workshop on Application Specific Processors, Dec. 2003, San Diego, CA. cited by applicant

Wooster et al., “HTTPDUMP Network HTTP Packet Snooper”, Apr. 25, 1996. cited by applicant

Worboys, “GIS: A Computing Perspective”, 1995, pp. 245-247, 287, Taylor & Francis Ltd. cited by applicant

Yamaguchi et al., “High Speed Homology Search with FPGAs”, Proceedings Pacific Symposium on Biocomputing, Jan. 3-7, 2002, pp. 271-282, vol. 7, Online, Lihue, Hawaii, USA. cited by applicant

Yan et al., “Enhancing Collaborative Spam Detection with Bloom Filters”, 2006, IEEE, pp. 414-425. cited by applicant

Yoshitani et al., “Performance Evaluation of Parallel vol. Rendering Machine Re Volver/C40”, Study Report of Information Processing Society, Mar. 5, 1999, pp. 79-84, vol. 99, No. 21. cited by applicant

Ziv et al., “A Universal Algorithm for Sequential Data Compression”, IEEE Trans. Inform. Theory, IT-23(3): 337-343 (1997). cited by applicant

“Decision Granting Institution of Inter Partes Review” for U.S. Pat. No. 10,121,196; IPR2020-01006, dated Dec. 4, 2020, pp. 1-22. cited by applicant

Office Action for U.S. Appl. No. 17/903,236 dated Apr. 11, 2024. cited by applicant

Primary Examiner: Gart; Matthew S

Assistant Examiner: Nguyen; Liz P

Attorney, Agent or Firm: Thompson Coburn LLP

Background/Summary

CROSS-REFERENCE AND PRIORITY CLAIM TO RELATED PATENT APPLICATIONS (1) This patent application is a continuation of U.S. patent application Ser. No. 16/180,377, filed Nov. 5, 2018, and entitled “Offload Processing of Data Packets Containing Financial Market Data”, now U.S. Pat. No. 10,963,962, which is a divisional of U.S. patent application Ser. No. 13/833,098, filed Mar. 15, 2013, and entitled “Offload Processing of Data Packets Containing Financial Market Data”, now U.S. Pat. No. 10,121,196, which claims priority to U.S. provisional patent application Ser. No. 61/616,181, filed Mar. 27, 2012, and entitled “Offload Processing of Data Packets Containing Financial Market Data”, the entire disclosures of each of which are incorporated herein by reference.

INTRODUCTION

(1) The process of trading financial instruments may be viewed broadly as proceeding through a cycle as shown in FIG. 1. At the top of the cycle is the exchange which is responsible for matching up offers to buy and sell financial instruments. Exchanges disseminate market information, such as the appearance of new buy/sell offers and trade transactions, as streams of events known as market data feeds. Trading firms receive market data from the various exchanges upon which they trade. Note that many traders manage diverse portfolios of instruments requiring them to monitor the state of multiple exchanges. Utilizing the data received from the exchange feeds, trading systems make trading decisions and issue buy/sell orders to the financial exchanges. Orders flow into the exchange where they are inserted into a sorted “book” of orders, triggering the publication of one or more events on the market data feeds.

(2) In an attempt to promptly deliver financial information to interested parties such as traders, a variety of electronic trading platforms have been developed for the purpose of ostensible “real time” delivery of streaming bid, offer, and trade information for financial instruments to traders. FIG. 2 illustrates an exemplary platform that is currently known in the art. As shown in FIG. 2, the electronic trading platform **200** comprises a plurality of functional units **202** that are configured to carry out data processing operations such as the ones depicted in units **202**, whereby traders at workstations **204** have access to financial data of interest and whereby trade information can be sent to various exchanges or other outside systems via output path **210**. The purpose and details of the functions performed by functional units **202** are well-known in the art. A stream **206** of financial data arrives at the system **200** from an external source such as the exchanges themselves (e.g., NYSE, NASDAQ, etc.) over private data communication lines or from extranet providers such as Savvis or BT Radians. The financial data source stream **206** comprises a series of messages that individually represent a new offer to buy or sell a financial instrument, an indication of a completed sale of a financial instrument, notifications of corrections to previously-reported sales of a financial instrument, administrative messages related to such transactions, and the like. As used herein, a “financial instrument” refers to a contract representing equity ownership, debt or credit, typically in relation to a corporate or governmental entity, wherein the contract is saleable. Examples of “financial instruments” include stocks, bonds, commodities, currency traded on currency markets, etc. but would not include cash or checks in the sense of how those items are used outside financial trading markets (i.e., the purchase of groceries at a grocery store using cash or check would not be covered by the term “financial instrument” as used herein; similarly, the withdrawal of \$100 in cash from an Automatic Teller Machine using a debit card would not be covered by the term “financial instrument” as used herein). Functional units **202** of the system then operate on stream **206** or data derived therefrom to carry out a variety of financial processing tasks. As used herein, the term “financial market data” refers to the data contained in or derived from a series of messages that individually represent a new offer to buy or sell a financial instrument, an indication of a completed sale of a financial instrument, notifications of corrections to previously-reported sales of a financial instrument, administrative messages related to such transactions, and the like. The term “financial market source data” refers to a feed of financial market data directly from a data source such as an exchange itself or a third party provider (e.g., a Savvis or BT Radianz provider). The term “financial market secondary data” refers to financial market data that has been derived from financial market source data, such as data produced by a feed compression operation, a feed handling operation, an option pricing

operation, etc.

(3) Financial data applications require fast access to large volumes of financial market data, and latency is an ever present technical problem in need of ever evolving solutions in the field of processing financial market data. As depicted in FIG. 2, the consumption, normalization, aggregation, and distribution of financial market data are key elements in a system that processes financial market data. For a broad spectrum of applications, platform architects seek to minimize the latency of market data processing and distribution, while minimizing the space and power required to host the market data processing and distribution elements. As described in the following patents and patent application, significant performance, efficiency, and scalability improvements can be achieved by leveraging reconfigurable hardware devices and other types of co-processors to integrate and consolidate market data consumption, normalization, aggregation, enrichment, and distribution functions: U.S. Pat. Nos. 7,840,482, 7,921,046, and 7,954,114 as well as the following published patent applications: U.S. Pat. App. Pub. 2007/0174841, U.S. Pat. App. Pub. 2007/0294157, U.S. Pat. App. Pub. 2008/0243675, U.S. Pat. App. Pub. 2009/0182683, U.S. Pat. App. Pub. 2009/0287628, U.S. Pat. App. Pub. 2011/0040701, U.S. Pat. App. Pub. 2011/0178911, U.S. Pat. App. Pub. 2011/0178912, U.S. Pat. App. Pub. 2011/0178917, U.S. Pat. App. Pub. 2011/0178918, U.S. Pat. App. Pub. 2011/0178919, U.S. Pat. App. Pub. 2011/0178957, U.S. Pat. App. Pub. 2011/0179050, U.S. Pat. App. Pub. 2011/0184844, WO Pub. WO 2010/077829, U.S. Pat. App. Pub. 2012/0246052, and U.S. Pat. App. Ser. No. 61/570,670, entitled “Method and Apparatus for Low Latency Data Distribution”, filed Dec. 14, 2011, the entire disclosures of each of which are incorporated herein by reference. These concepts can be extended to various market data processing tasks as described in the above-referenced and incorporated patents and patent applications. Similarly, the above-referenced and incorporated Pat. App. Ser. No. 61/570,670 demonstrates how the systems responsible for the distribution of real-time financial data can be greatly enhanced via the use of novel communication protocols implemented in reconfigurable hardware devices and other types of co-processors.

(4) In accordance with various embodiments disclosed herein, the inventors further disclose various methods, apparatuses, and systems for offloading the processing of data packets that contain financial market data. In exemplary embodiments, various processing tasks are offloaded from an electronic trading platform to one or more processors upstream or downstream from the electronic trading platform. It should be understood that the term upstream in this context is meant to identify a directional flow with respect to data that is moving to an electronic trading platform, in which case an offload processor upstream from the electronic trading platform would process financial market data flowing toward the electronic trading platform. Similarly, in this context downstream is meant to identify a directional flow with respect to data that is moving away from an electronic trading platform, in which case an offload processor downstream from the electronic trading platform would process financial market data flowing out of the electronic trading platform.

(5) In some embodiments, the offloaded processing can be moved into the data distribution network for financial market data. For example, one or more of the offloaded financial market data processing tasks described herein can be implemented in one or more network elements of the data distribution network, such as a switch within the data distribution network. Disclosed herein are exemplary embodiments where a number of market data consumption, normalization, aggregation, enrichment, and distribution functions can be embedded within the elements that comprise the market data feed network **214**. Conceptually, these embodiments offload processing tasks typically performed by downstream processing elements **202** such as feed handlers and virtual order books. The inventors also disclose a number of market data distribution functions that can be embedded within the network elements that comprise the financial application data network **208**. Conceptually, these embodiments effectively offload processing tasks typically performed by ticker plants, messaging middleware, and downstream applications. Offloading these tasks from traditional platform components and embedding them in network elements may obviate some platform components, improve the performance of some components, reduce the total amount of space and power required by the platform, achieve higher system throughput, and deliver lower latency market data to consuming applications.

(6) These and other features and advantages of the present invention will be apparent to those having ordinary skill in the art upon review of the teachings in the following description and drawings.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 illustrates an exemplary process cycle for trading financial instruments.
- (2) FIG. 2 illustrates an exemplary electronic trading platform.
- (3) FIGS. 3-6 illustrate exemplary embodiments for offload processors that provide repackaging functionality.
- (4) FIG. 7 illustrates an exemplary system where an offload processor is deployed upstream from one or more electronic trading platform(s).
- (5) FIG. 8 illustrates an exemplary system where an intelligent feed switch is positioned within the market data feed network of an electronic trading platform.
- (6) FIG. 9 illustrates an exemplary system where conventional switches are used to aggregate financial market data feeds for delivery to an intelligent feed switch.
- (7) FIG. 10 illustrates an exemplary system where conventional switches are used to aggregate financial market data feeds for delivery to multiple intelligent feed switches.
- (8) FIG. 11 depicts an exemplary electronic trading platform with an intelligent feed switch deployed in the market data network.
- (9) FIG. 12 illustrates the system of FIG. 11 including a logical diagram of functions performed by a typical feed handler in an electronic trading platform.
- (10) FIG. 13 illustrates the system of FIG. 11 but where several functions are offloaded from the feed handler to the intelligent feed switch.
- (11) FIG. 14 illustrates an exemplary electronic trading platform that includes one or more ticker plant components.
- (12) FIG. 15 illustrates the system of FIG. 14 but where several functions are offloaded from a ticker plant to the intelligent feed switch.
- (13) FIG. 16 illustrates an exemplary system where latency-sensitive trading applications consume data directly from an intelligent feed switch.
- (14) FIG. 17 illustrates an example of redundant feed arbitration.
- (15) FIG. 18 illustrates an example of a line arbitration offload engine.
- (16) FIG. 19 illustrates an example of a packet mapping offload engine.
- (17) FIG. 20 illustrates an exemplary processing module configured to perform symbol-routing and repackaging.
- (18) FIG. 21 illustrates an exemplary intelligent feed switch that provides multiple ports of 10 Gigabit Ethernet connectivity.
- (19) FIG. 22 illustrates an exemplary intelligent feed switch wherein the switch device is replaced by another FPGA device with a dedicated memory cache.
- (20) FIG. 23 illustrates an exemplary intelligent feed switch wherein a single FPGA device is utilized.
- (21) FIG. 24 illustrates an exemplary intelligent distribution switch positioned downstream of market data normalization components in an electronic trading platform.
- (22) FIG. 25 illustrates an exemplary intelligent distribution switch that hosts one or more distribution functions.

DETAILED DESCRIPTION

(23) A. Offload Processor:

(24) Thus, in an exemplary embodiment, the inventors disclose that an offload processor can be configured to process incoming data packets, where each of at least a plurality of the incoming data packets contain a plurality of financial market data messages, and wherein the financial market data messages comprise a plurality of data fields describing financial market data for a plurality of financial instruments. Thus, the payload of each incoming data packet can comprise one or more financial market data messages. Such an offload processor can filter and repackage the financial market data into

outgoing data packets that are grouped into outgoing data packets is grouped using a criterion different than the criterion upon which financial market data was grouped into the incoming data packets. This permits the offload processor to serve a valuable role in generating a new set of customized outgoing data packets from incoming data packets. In various exemplary embodiments of such an offload processor, the offload processor can alleviate the processing burden on the downstream electronic trading platform(s).

(25) Examples of such an offload processor are shown in FIGS. 3-6. FIG. 3 depicts an exemplary offload processor **300** that is configured to receive as an input a consolidated (i.e., multiplexed) stream of incoming data packets from different financial markets. As shown in FIG. 3, each incoming data packet has a payload that contains multiple financial market data messages from the same financial market. Thus, a plurality of financial market data messages from the feed for Financial Market 1 (e.g., NYSE) are combined in the same packet (e.g., where financial market data message FMDM1 (Mkt 1) is a new offer to buy stock for Company A from the NYSE, FMDM2 (Mkt 1) is a new offer to sell stock for Company B from the NYSE, and where FMDM3 (Mkt 1) is a notification of a completed trade on stock for Company C from the NYSE), while a plurality of financial market data messages from the feed for Financial Market 2 (e.g., NASDAQ) are combined in the same packet, and so on. The offload processor **300** performs financial market data filtering and repackaging between incoming and outgoing data packets such that the outgoing financial market data packets contain financial market data messages that are organized using a different criterion. Thus, the offload processor filters and sorts the financial market data from the different markets by a criterion such as which downstream data consumers have expressed an interest in such financial market data. In this fashion, the offload processor **300** can mix payload portions of incoming data packets on a criterion-specific basis to generate outgoing data packets with newly organized payloads. For example, data consumer A may have an interest in all new messages relating a particular set of financial instruments (e.g., IBM stock, Apple stock, etc.) regardless of which market served as the source of the messages on such instruments. Another data consumer, Consumer B, may have similar interests in a different set of financial instruments. In such a case, the offload processor can be configured to re-group the financial market data into the outgoing data packets around the interests of particular downstream consumers. Thus, FIG. 3 also shows outgoing data packets that are consumer-specific. As can be seen, the payloads of these consumer-specific data packets comprise financial market data messages from different markets that arrived in different incoming data packets.

(26) Exemplary processing pipelines that can be employed by the offload processor to provide such sorting and repackaging functions are described below in connection with FIGS. 13, 15, and 20.

(27) In another exemplary embodiment, an offload processor can be configured to perform packet mapping functions on incoming data packets from various financial market data feeds.

(28) FIG. 4 depicts another exemplary embodiment of an offload processor **300** that provides repackaging functionality. In the example of FIG. 4, the offload processor receives a plurality of streams of incoming data packets, where each stream may be market-specific (e.g., an input stream of data packets from the NYSE on a first port and an input stream of data packets from NASDAQ on a second port). The offload processor **300** of FIG. 4 can then repackage the financial market data in these incoming data packets into outgoing data packets as previously discussed.

(29) FIG. 5 depicts another exemplary embodiment of an offload processor **300** that provides repackaging functionality. In the example of FIG. 5, the offload processor produces multiple output streams of outgoing data packets, where each output stream may be criterion-specific (e.g., an output stream of data packets destined for Consumer A from a first port and an output stream of data packets destined for Consumer B from a second port, and so on). The stream of incoming data packets can be a consolidated stream as described in connection with FIG. 3.

(30) FIG. 6 depicts another exemplary embodiment of an offload processor **300** that provides repackaging functionality. In the example of FIG. 6, the offload processor produces multiple output streams of outgoing data packets from multiple input streams of incoming data packets, where the input streams can be like those shown in FIG. 4 while the output streams can be like those shown in FIG. 5.

(31) The offload processor **300** can take any of a number of forms, including one or more general purpose processors (GPPs), reconfigurable logic devices (such as field programmable gate arrays (FPGAs), application-specific integrated circuits (ASICs), graphics processing units (GPUs), and chip

multiprocessors (CMPs), as well as combinations thereof.

(32) As used herein, the term “general-purpose processor” (or GPP) refers to a hardware device having a fixed form and whose functionality is variable, wherein this variable functionality is defined by fetching instructions and executing those instructions, of which a conventional central processing unit (CPU) is a common example. Exemplary embodiments of GPPs include an Intel Xeon processor and an AMD Opteron processor. As used herein, the term “reconfigurable logic” refers to any logic technology whose form and function can be significantly altered (i.e., reconfigured) in the field post-manufacture. This is to be contrasted with a GPP, whose function can change post-manufacture, but whose form is fixed at manufacture. Furthermore, as used herein, the term “software” refers to data processing functionality that is deployed on a GPP or other processing devices, wherein software cannot be used to change or define the form of the device on which it is loaded, while the term “firmware”, as used herein, refers to data processing functionality that is deployed on reconfigurable logic or other processing devices, wherein firmware may be used to change or define the form of the device on which it is loaded.

(33) Thus, in embodiments where the offload processor **300** comprises a reconfigurable logic device such as an FPGA, hardware logic will be present on the device that permits fine-grained parallelism with respect to the different operations that the offload processor performs, thereby providing the offload processor with the ability to operate at hardware processing speeds that are orders of magnitude faster than would be possible through software execution on a GPP. Moreover, by leveraging such fine-grained parallelism, processing tasks can be intelligently engineered into processing pipelines deployed as firmware in the hardware logic on the FPGA. With such a pipeline, downstream pipeline modules can perform a processing task on data that was previously processed by upstream pipelined modules while the upstream pipeline modules are simultaneously performing other processing tasks on new data, thereby providing tremendous throughput gains. Furthermore, other types of offload processors that provide parallelized processing capabilities can also contribute to improved latency and throughput.

(34) FIG. 7 depicts an exemplary system where the offload processor **300** is deployed upstream from one or more electronic trading platform(s) (ETP(s)) **700**. Each ETP **700** may include one or more data consumers within it, and the outgoing data packets from the offload processor **300** can be customized to each consumer.

(35) Furthermore, in additional exemplary embodiments, the offload processor can perform other functions in addition to or instead of the repackaging operations illustrated by FIGS. 3-6. For example, the offload processor can be configured to perform packet mapping as described below in connection with FIG. 19.

(36) As noted, when positioned upstream from an electronic trading platform, the offload processor can be employed in a network element resident in a data distribution network for financial market data. Examples of network elements include repeaters, switches, routers, and firewalls. A repeater embodiment, a single input port and single output port device, may be viewed as a “smart” link where data is processed as it flows through the network link. In a preferred embodiment, such a network element can be a network switch. As such, the inventors disclose various embodiments of a network switch that offloads various processing tasks from electronic trading platforms, including embodiments of an intelligent feed switch and embodiments of an intelligent distribution switch, as described below.

(37) B. Intelligent Feed Switch:

(38) A common practice in financial exchange and electronic trading platform architecture is to achieve greater scale by “striping the data” across multiple instances of the platform components responsible for data transmission, consumption, and processing. If the data is imagined to flow vertically through a depiction of the overall system, then this approach to scale is often termed “horizontal scaling”. This approach is accepted in the industry as the most viable approach from an overall platform perspective, as the escalating rate of market data messages (doubling every 6 to 11 months) is outpacing the technology improvements available to individual components in the platform.

(39) In order to facilitate data striping, some feed sources (typically exchanges) divide a market data feed into multiple “lines” where a given line carries a proper subset of the market data published by the financial exchange. Typically, all of the market data updates associated with a given financial instrument is transmitted on a single line. The assignment of a given financial instrument to a line may be static or dynamic. Static assignments typically partition the set of instruments by using the starting characters in

an instrument symbol and assigning an alphabet range to a given line. For example, consider a feed partitioned into four lines. Line 0 carries updates for financial instruments whose symbol begins with letters “A” through “F”; line 1 carries updates for symbols beginning with letters “G” through “M”; line 2 carries updates for symbols beginning with letters “N” through “S”; line 3 carries updates for symbols beginning with letters “T” through “Z”. Dynamic line assignments are typically performed as follows. A static mapping line transmits information to feed consumers communicating the number of data lines, the address(es) of the data lines, and the mapping of financial instruments to each data line.

(40) Similarly, financial exchanges typically enforce striping across the ports provided for order entry. A financial exchange provides multiple communication ports to which market participants establish connections and enter orders to electronically buy and sell financial instruments. Exchanges define the subset of financial instruments for which orders are accepted on a given port. Typically, exchanges statically define the subset of financial instruments by using the starting character(s) in the instrument symbol. They assign an alphabet range to a given port. For example, consider an exchange that provides four ports to a given participant. Port 0 accepts orders for financial instruments whose symbol begins with letters “A” through “F”; port 1 accepts orders for symbols beginning with letters “G” through “M”; port 2 accepts orders for symbols beginning with letters “N” through “S”; port 3 accepts orders for symbols beginning with letters “T” through “Z”.

(41) The striping of data by exchanges, across multiple market data feed lines as well as multiple order entry ports, dictates a horizontally scaled architecture for electronic trading platforms. Trading applications are typically responsible for trading a subset of the financial instruments. Each application consumes the market data updates associated with its subset of financial instruments and generate orders for those instruments. Implementing a horizontally scaled system is straightforward for a platform that receives data from and transmits orders to a single market. The design task is significantly complicated when the trading platform receives data from multiple exchanges, computes pan-market views of financial instruments, and transmits orders to multiple exchanges.

(42) Each market data feed source implements its own striping strategy. Note that some market data feeds are not striped at all and employ a single line. The subsets of financial instruments associated with the lines on one market data feed may be different from the subsets of financial instruments associated with the lines on another market data feed. Therefore, the updates associated with financial instruments processed by a given component can be sourced from different sets of lines from each market data feed. These factors significantly complicate the market data processing and distribution components that are responsible for delivering normalized market data to downstream applications, especially when composite, pan-market views of financial instruments are required.

(43) Disclosed herein are multiple variants of an Intelligent Feed Switch (IFS) that offloads numerous market data consumption, normalization, aggregation, enrichment, and distribution functions from downstream components such as feed handlers, virtual order books, or more generally, ticker plants. The specific functions performed by variants of the IFS are described in the sections below. As previously mentioned, utilizing an IFS in the market data feed network provides performance, efficiency, functionality, and scalability benefits to electronic trading platforms.

(44) 1. IFS Architecture:

(45) The IFS can be implemented on a wide variety of platforms that provide the necessary processing and memory resources, switching resources, and multiple physical network ports. Just as network switches can be built at various scales, two ports up to thousands of ports, the IFS can be scaled to meet the needs of electronic trading platforms of varying scale. In the embodiment shown in FIG. 21, the IFS provides multiple ports of 10 Gigabit Ethernet connectivity, in addition to a 10/100/1000 Ethernet port for management and control. An FPGA that is resident within the switch can provide fine-grained parallel processing resources for offload engines as previously noted. The memory cache provides dedicated high-speed memory resources for the offload engines resident on the FPGA. The memory cache may be implemented in Synchronous Dynamic Random Access Memory (SDRAM), Synchronous Random Access Memory (SRAM), a combination of the two, or other known memory technologies. A dedicated Ethernet switch ASIC increases the port count of the IFS using existing, commodity switching devices and allows traffic to bypass the offload engines in the FPGA. The FPGA is directly connected to the switching device by consuming one or more ports on the switching device. The amount of

communication bandwidth between the FPGA and switching device can be scaled by increasing the number of ports dedicated to the interface. The FPGA may also provide one or more ports for external connectivity, adding to the total number of ports available on the IFS. In addition to providing standard protocol connectivity, e.g. Ethernet, the ports that are directly connected to the FPGA can be leveraged to implement custom protocols. For example, if multiple Intelligent Feed Switches are interconnected, the FPGAs inside the switches may implement a custom protocol that eliminates unnecessary overhead. Similarly, if a custom Network Interface Card (NIC) containing an FPGA directly connected to the physical network port(s) is used in a server connected to the IFS, a custom protocol can be employed between the IFS and the server. The control processor provides general purpose processing resources to control software. A standard operating system (OS) such as Linux is installed on the control processor. Configuration, control, and monitoring software interfaces with the FPGA device via a standard system bus, preferably PCI Express. The control processor also features a system bus interface to the switch device.

(46) FIG. 22 shows another embodiment of the IFS wherein the switch device is replaced by another FPGA device with a dedicated memory cache. Note that the peer-to-peer (P2P) interface between the FPGA devices need not utilize a standard network protocol, such as Ethernet, but may use a low-overhead protocol for communicating over high speed device interconnects. This architecture increases the amount of processing resources available for offload functions and allows custom network protocols to be supported on any port. Also note that additional FPGAs can be interconnected to scale the number of external ports provided by the IFS.

(47) FIG. 23 shows another embodiment of the IFS wherein a single FPGA device is utilized. This architecture can minimize cost and complexity. The number of physical ports supported is subject to the capabilities of the selected FPGA device. Note that some devices include embedded general purpose processors capable of hosting configuration, control, and monitoring applications.

(48) Note that other processing resources such as chip multi-processors (CMPs), graphics processing units (GPUs), and network processing units (NPUs) may be used in lieu of an FPGA. An example of a network switch platform that may be suitable for use as an intelligent switch to process financial market data is the Arista Application Switch 7124FX from Arista Networks, Inc. of Santa Clara, CA.

(49) 2. Platform Architecture with IFS:

(50) As shown in FIG. 8, the IFS can be positioned within the market data feed network of the electronic trading platform. In some market data networks, a single IFS may be capable of providing the required number of switch ports, processing capacity, and data throughput. The number of switch ports required depends on the number of physical network links carrying input market data feeds and the number of physical network links connecting to downstream platform components. The amount of processing capacity required depends on the tasks performed by the IFS and the requirements imposed by the input market data feeds. The data throughput depends on the aggregate data rates of input market data feeds and aggregate data rates of output streams delivered to platform components.

(51) If the aforementioned requirements exceed the capacity of a single IFS, then a multi-element network can be constructed that includes the IFS. As shown in FIG. 9, multiple conventional switch elements can be used to aggregate the data from the physical network links carrying market data feeds. For example, a conventional switch could be used to aggregate data from forty (40) 1 Gigabit Ethernet links into four (4) 10 Gigabit Ethernet links for transfer to the IFS. This reduces the number of upstream ports required by the IFS. As shown in FIG. 10, multiple Intelligent Feed Switches can be used if the requirements exceed the capacity of a single IFS. In this example, multiple IFS elements consume aggregated data from upstream conventional switches, then distribute data to downstream platform elements. The network architectures in FIGS. 9 and 10 are exemplary but not exhaustive. The IFS can be combined with other switch elements to form large networks, as is well-known in the art.

(52) FIG. 11 presents a simplified diagram of a conventional electronic trading platform with an IFS deployed in the market data network. In this arrangement, the IFS offloads one or more functions from the downstream feed handler components. FIG. 12 provides a logical diagram of the functions performed by a typical feed handler in a conventional electronic trading platform. A description of the specific functions and how they can be offloaded to the IFS are described in detail in the sections below. FIG. 13 provides a logical diagram of a conventional electronic trading platform with numerous feed handler

function performed by the IFS. Note that the only remaining functions performed by the feed handler components are message parsing, business logic and message normalization, and subscription-based distribution. Note that we later describe an embodiment capable of further offloading the feed handler components from subscription-based distribution. With a substantially reduced workload, feed handler components can be re-engineered to be more simple, efficient, and performant. As a result the number of discrete feed handler components required by the electronic trading platform can be substantially reduced. The latency associated with market data normalization and distribution can be substantially reduced, resulting in advantages for latency-sensitive trading applications. Furthermore, the amount of space and power required to host the electronic trading platform can be substantially reduced, resulting in simplified system monitoring and maintenance as well as reduced cost.

(53) FIG. 14 presents a simplified diagram of an electronic trading platform that includes one or more ticker plant components that integrate multiple components in the conventional electronic trading platform. An example of an integrated ticker plant component that leverages hardware acceleration and offload engines is described in the above-referenced and incorporated patents and patent applications (see, for example, U.S. Pat. No. 7,921,046, U.S. Pat. App. Pub. 2009/0182683, and WO Pub. WO 2010/077829). Even integrated ticker plant components such as these can benefit from offloading functions to an IFS. As shown in FIG. 15, the IFS can offload the feed handling tasks reflected in FIG. 13, as well as additional functions such as price aggregation, event caching, top-of-book quote generation, and data quality monitoring. A description of these functions and how they can be offloaded to an IFS is provided in subsequent sections. Offloading these functions can boost the capacity of an integrated ticker plant component, reducing the need to horizontally scale. An IFS can also simplify the task of horizontally scaling with multiple integrated ticker plant components. For example, consider a platform architecture where three ticker plant components are used and horizontal scaling is achieved by striping the symbol range across the ticker plant components. The first ticker plant is responsible for processing updates for instrument symbols beginning with characters “A” through “H”. The IFS is capable of ensuring that the first ticker plant only receives updates for the assigned set of instruments by performing the symbol routing and repackaging functions depicted in FIG. 15. Note that other functions predicate the symbol routing function as described subsequently. Striping the data in this way allows each ticker plant component to retain the ability to compute composite, or pan-market, views of financial instruments. Examples of hardware-accelerated processing modules for computing composite quote and order book views are described in the above-referenced and incorporated U.S. Pat. No. 7,921,046 and WO Pub. WO 2010/077829.

(54) Some latency-sensitive trading applications require minimal data normalization in order to drive their trading strategies. Some of these applications may be able to directly consume data from an IFS, as shown in FIG. 16. This eliminates additional network hops and processing from the datapath, thus reducing the latency of the data delivered to the applications. This latency reduction can provide advantages to these latency-sensitive trading applications.

(55) As shown in FIG. 8, the IFS is positioned within the market data feed network, and represents the physical embodiment of that network.

(56) 3. Packet Mapping:

(57) As shown in FIGS. 13 and 15, the IFS may be configured to offload one or more functions from downstream feed consumers. The same set of functions may not be performed for every feed flowing through the IFS. Furthermore, the way in which each function is performed may vary by feed, as feed sources employ different message formats, field identifiers, datatypes, compression schemes, packet formats, transmission protocols, etc. In order to correctly perform the prescribed functions on a given packet, the IFS must first identify the feed to which a given packet belongs, then retrieve the necessary information about how packets belonging to the given feed are to be handled. In order to do so, the IFS preferably maintains a mapping table using a tuple such as the IP<source address, destination address, protocol> tuple to identify the feed to which a packet belongs (additional optional members of the tuple may include a source port number, a destination port number, and a transport protocol port number). Preferably, the embedded processor in the IFS utilizes a hash table, where the <source address, destination address, protocol> tuple is used as input to the hash function. Preferably, a control processor in the IFS configures the hash function and maintains the hash table. At minimum in this example, the

entry in the table contains a feed identifier. The additional information about how packets belonging to the feed should be handled may be stored directly in the hash table, or in a separate table indexed by the feed identifier. The additional information may include one or more of the following pieces of meta-data: Market identification code (MIC); a unique identifier for the exchange/market. Preferably, this code would be a binary enumeration of the ISO 10383 market identification codes (MIC) for the markets supported by the IFS. For example, XNYS is the MIC for the New York Stock Exchange which may be assigned an enumerated value in order to consume minimal space in the meta-data table and pre-normalized messages. Data source identification code (DSIC); a unique identifier for the specific feed. Note that multiple feeds may carry market updates for the same market. For example, updates for equities traded on the NYSE are reported by multiple feeds: the Consolidated Quote System (CQS), Consolidated Tape System (CTS), NYSE Quotes, NYSE Trades, NYSE OpenBook Ultra, etc. Each feed, or data source, is assigned a unique tag. Similar to the market codes, the data source codes are assigned an enumerated value in order to consume minimal space in the meta-data table and pre-normalized messages. Line identification code (LIC); a unique identifier for the specific line within the feed. Similar to the MIC and DSIC, each unique line is assigned a unique tag. The line identifiers configured on the IFS are preferably assigned an enumerated value in order to consume minimal space in the meta-data table and pre-normalized messages. A flag indicating if the feed utilizes FIX/FAST encoding FAST decoding templates (if necessary), or template specifying how to parse the packet into messages. FIX decoding templates, or template specifying how to parse messages into fields. Template specifying field datatype conversions to perform. Field identifiers and/or offsets for fields comprising the instrument symbol. Field identifier or offset for message sequence number (if necessary)

(58) This meta-information can be propagated to downstream offload engines in the IFS, along with the packet, as shown in FIG. 19. The configuration, control, and table management logic configures the hash function and table entries. This logic is preferable hosted on a co-resident control processor, preferably as a pipelined processing engine.

(59) 4. Redundant Feed Arbitration:

(60) In order to allow a market data feed to be routed across multiple networks, the Internet Protocol (IP) is ubiquitously used as the network protocol for market data feed distribution. Feed sources typically employ one of two transport protocols: Transmission Control Protocol (TCP) or Unreliable Datagram Protocol (UDP). TCP provides a reliable point-to-point connection between the feed source and the feed consumer. Feed consumers initiate a connection with the feed source, and the feed source must transmit a copy of all market data updates to each feed consumer. Usage of TCP places a large data replication load on the feed source, therefore it is typically used for lower bandwidth feeds and/or feeds with a restricted set of consumers. UDP does not provide reliable transmission, but does include multicast capability. Multicast allows the sender to transmit a single copy of a datagram to multiple consumers. Multicast leverages network elements to perform the necessary datagram replication. An additional protocol allows multicast consumers to “join” a multicast “group” by specifying the multicast address assigned to the “group”. The sender sends a single datagram to the group address and intermediary network elements replicate the datagram as necessary in order to pass a copy of the datagram to the output ports associated with consumers that have joined the multicast group.

(61) While providing for efficient data distribution, UDP multicast is not reliable. Datagrams can be lost in transit for a number of reasons: congestion within a network element causes the datagram to be dropped, a fault in a network link corrupts one or more datagrams transiting the link, etc. While there have been numerous reliable multicast protocols proposed from academia and industry, none have found widespread adoption. Most market data feed sources that utilize UDP multicast transmit redundant copies of the feed, an “A side” and a “B side”. Note that more than two copies are possible. For each “line” of the feed, there is a dedicated multicast group, an “A” multicast group and a “B” multicast group. Typically, the feed source ensures that each copy of the feed is transmitted by independent systems, and feed consumers ensure that each copy of the feed transits an independent network path. Feed consumers then perform arbitration to recover from data loss on one of the redundant copies of the feed.

(62) Note that a packet may contain one or more market data update messages for one or more financial instruments. Typically, feed sources assign a monotonically increasing sequence number to each packet

transmitted on a given “line”. This simplifies the task of detecting data loss on a given line. If the most recently received packet contains a sequence number of **5893**, then the sequence number of the next packet should be **5894**. When using redundant UDP multicast groups, feed sources typically transmit identical packets on the redundant multicast groups associated with a line. For example, packet sequence number **3839** on the A and B side of the feed contains the same market data update messages in the same order. This simplifies the arbitration process for feed consumers.

(63) FIG. **17** provides a simple example of redundant feed arbitration. The sequence of packets for a single pair of redundant lines is shown. Time progresses vertically, with packet **5894** received first from line 1A, packet **5895** received second from line 1A, etc. A line arbiter forwards the packet with the next sequence number, regardless of which “side” the packet arrives on. When the redundant copy of the packet is received on the other side, it is dropped. As depicted in FIG. **17**, one of the redundant sides typically delivers a packet consistently prior to the other side. If the arbiter receives a packet with a sequence number greater than the expected sequence number, it detects a gap on one of the redundant lines. The arbiter can be configured to wait a configured hold time to see if the missing packet is delivered by the other side. The difference between the arrival times of copies of the same packet on the redundant lines is referred to as the line skew. In order to be effective, the hold time can be configured to be greater than the average line skew. If the missing packet does arrive on the redundant side prior to the expiration of the hold time, then a gap is registered for the particular feed line.

(64) When line gaps occur there are a number of recovery and mitigation strategies that can be employed. The arbiter typically reports the missing sequence numbers to a separate component that manages gap mitigation and recovery. If the feed provides retransmission capabilities, then the arbiter may buffer packets on both sides until the missing packets are returned by the gap recovery component.

(65) Some feeds sequence updates on a per-message basis or a per-message/per-instrument basis. In these cases, a packet sequence number may not be monotonically increasing or may not be present at all. Typically, arbitration is performed among one or more copies of a UDP multicast feed; however, arbitration can occur among copies of the feed delivered via different transmission protocols (UDP, TCP, etc.). In these scenarios, the content of packets on the redundant copies of the feed may not be identical. The transmitter of packets on the A side may packetize the sequence of market data update messages differently from the transmitter on the B side. This requires the IFS to parse packets prior to performing the arbitration function.

(66) The line identification code (LIC) provided in the meta-data associated with the packet allows the IFS to perform the appropriate line arbitration actions for a given packet. If the packet belongs to an unarbitrated TCP flow, then the packet may bypass the line arbitration and gap detection engine. If the line requires dictates arbitration at the message-level as opposed to the packet level, then the IFS first routes the packet to parsing and decoding engines. The line arbitration and gap detection function may be performed by multiple parallel engines. The LIC may also be used to route the packet to the appropriate engine handling arbitration for the associated feed line. Furthermore, the LIC is used to identify the appropriate arbitration buffer into which the packet should be inserted.

(67) FIG. **18** provides an example of a line arbitration offload engine, which is preferably implemented in a pipelined processing engine. For each input line, the arbiter maintains a packet buffer to store the packets received from the redundant sides of the feed line. The example in FIG. **18** demonstrates two-arbitration; additional buffers are provisioned if multi-way arbitration is performed. For feeds transmitted via UDP, it is possible for packets on a given multicast group to be delivered in out-of-sequence, if the packets traverse different paths through the network. The packet buffers in the arbiter may optionally provide for resequencing by inserting each new packet in the proper sequence in the buffer. Typically market data networks are carefully designed to minimize latency and tightly control routing, thus out-of-sequence delivery is typically not a problem. Thus, arbiter functions typically omit resequencing to reduce overhead and complexity.

(68) The compare, select and drop logic in the arbiter performs the core arbitration function as previously described. A register is used to maintain the next expected sequence number. The logic compares the sequence number of the packet residing at the head of each packet buffer. If a matching sequence number is found, the packet is forwarded. If the sequence number is less than the expected sequence number, the packet is dropped. If the sequence number is greater than the expected sequence number, the other buffer

or buffers are examined for the required packet. Note that this may require that multiple packets be read until a match is found, the buffer is empty, or a gap is detected. If a gap is detected the gap detection and reporting logic resets then starts the wait timer. If the expected packet sequence number does not arrive before the wait timer exceeds the value in the max hold time register, then a gap is reported to the gap mitigation and recovery engine with the missing packet sequence number range. Note that the gap detection and reporting logic may also report gap information to a control processor or to downstream monitoring applications via generated monitoring messages. If the gap mitigation and recovery engine is configured to request retransmissions, then the arbiter pauses until the gap mitigation and recovery engine passes the missing packet or packets to the arbiter or returns a retransmission timeout signal. The gap mitigation and recovery engine may be hosted on the same device as the arbiter, or it may be hosted on a control processor within the IFS.

(69) 5. Feed Pre-Normalization:

(70) In addition to performing line arbitration and gap detection, mitigation, and recovery, the IFS can perform one or more “pre-normalization” functions in order to simplify the task of downstream consumers. Following line arbitration, the IFS preferably decomposes packets into discrete messages. As previously described, feed sources typically pack multiple update messages in a single packet. Note that each feed may employ a different packetization strategy, therefore, the pre-normalization engine in the IFS utilizes the packet parsing templates retrieved by the packet mapping engine. Packet parsing techniques amenable to implementation in hardware and parallel processors are known in the art as described in the above-referenced and incorporated U.S. Pat. No. 7,921,046. If the feed associated with the packet utilizes FAST compression, then the pre-normalization engine must utilize the FAST decoding template in order to decompress and parse the packet into individual messages, as described in the above-referenced and incorporated U.S. Pat. No. 7,921,046.

(71) Once the packet is parsed into discrete messages, specific fields may be extracted from the messages in order to enable additional pre-normalization functions. Template-based parsing in offload engines is also addressed in the above-referenced and incorporated U.S. Pat. No. 7,921,046. Discrete messages and message fields are passed to downstream functions. Note that the message parsing engine may only extract specific fields required for downstream functions, as dictated by the templates included in the meta-data for the packet. For example, the parser may only extract the symbol field in order to enable symbol-based routing and repackaging. For some feeds, the symbol mapping function may require extraction of the order reference number in book update events. This can also be specified by the parsing template.

(72) Note that the message parsing logic can be configured to preserve the original structure of the message. Extracted fields, such as symbols and order reference numbers, can be added to the meta-data that accompanies the packet as it propagates through the IFS. By preserving the message structure, downstream consumer applications need not be changed when an IFS is introduced in the market data network. For example, an existing feed handler for the NASDAQ Total View feed need not change, as the format of the messages it processes still conforms to the feed specification. If the symbol-routing and repackaging function is applied, the existing feed handler will simply receive packets with messages associated with the symbol range for which it is responsible, but the message formats will conform to the exchange specification. This function is described in more detail below.

(73) The pre-normalization logic can also be configured to offload normalization logic from downstream consumers. For example, the parsing logic can be configured to perform FAST decompression and FIX parsing. Per the parsing templates in the meta-data, the fields in each message can be configured to a prescribed native data type. For example, an ASCII-encoded price field can be converted into a signed 32-bit integer, an ASCII-encoded string can be mapped to a binary index value, etc. The type-converted fields can then be aligned on byte or word boundaries in order to facilitate efficient consumption by consumers. The pre-normalization logic can maintain a table of downstream consumers capable of receiving the pre-normalized version of the feed. For example, the IFS may transmit pre-normalized messages on ports 3 through 8, but transmit the raw messages on ports 9 through 12.

(74) For some feeds, the IFS can be configured to append fields to the raw message, allowing consuming applications to be extended to leverage the additional fields to reap performance gains, without disrupting the function of existing consumers. For example, the IFS may append the MIC, DSIC, LIC,

and binary symbol index to the message. Additional appended fields may include, but are not limited to, message-based sequence numbers and high-resolution IFS transmit timestamps.

(75) As previously mentioned, the IFS can be configured to perform a symbol mapping function. The symbol mapping function assigns a binary symbol index to the financial instrument associated with the update event. This index provides a convenient way for downstream functions and consuming applications to perform processing on a per symbol basis. An efficient technique for mapping instrument symbols using parallel processing resources in offload engines is described in the above-referenced and incorporated U.S. Pat. No. 7,921,046. Note that some feeds provide updates on a per-order basis and some update events do not contain the instrument symbol, but only an order reference number. Feed consumers must map the order reference number to the associated symbol. An efficient technique for mapping order reference numbers to the mapped symbol index using parallel processing resources in offload engines is described in the above-referenced and incorporated WO Pub. WO 2010/077829. In order to perform the symbol mapping function, the computational resources in the IFS can include dedicated high-speed memory interfaces.

(76) As part of the pre-normalization function, the IFS may also assign one or more high-precision timestamps. For example, a timestamp may be assigned when the IFS receives a packet, a timestamp may be assigned immediately prior to transmitting a packet, etc. The high-precision timestamp preferably provides nanosecond resolution. In order to provide synchronized timestamps with downstream consumers, the time source used to assign the timestamps should be disciplined with a high-precision time synchronization protocol. Example protocols include the Network Time Protocol (NTP) and the Precision Time Protocol (PTP). The protocol engine can be co-resident with the offload engines in the IFS, but is preferably implemented in a control processor that disciplines a timer in the offload engines. As part of the pre-normalization function, the IFS may also assign additional sequence numbers. For example, the IFS may assign a per-message, per-symbol sequence number. This would provide a monotonically increasing sequence number for each instrument. These additional timestamps and sequence numbers may be appended to raw message formats or included in the pre-normalized message format, as described above.

(77) 6. Symbol-Based Routing and Repackaging:

(78) The symbol-based routing allows the IFS to deliver updates for a prescribed set of symbols to downstream components in the electronic trading platform. As shown in FIG. 16, the IFS can act as a subscription based routing and filtering engine for latency-sensitive applications that consume the raw or pre-normalized updates directly from the IFS. Similarly, the IFS can facilitate a horizontal scaling strategy by striping the incoming raw feed data by symbol within the market data feed network itself. This allows the IFS to deliver the updates for the prescribed symbol range to downstream feed handler or ticker plant components, without having to rely on additional processing capabilities in those components to perform this function. This can dramatically reduce data delivery latency and increase the processing capacity of those components.

(79) FIG. 20 depicts an exemplary processing module configured to perform symbol-routing and repackaging. Such a module is preferably implemented as a pipelined processing engine. As shown in FIG. 20, the symbol-routing and repackaging function first utilizes the symbol index to lookup an interest list in the interest list table. Note that additional fields such as the market identification code (MIC) and data source identification code (DSIC) may be used in addition to the symbol index to lookup an interest list. Similar to the interest-based filtering and replication discussed in the above-referenced and incorporated U.S. Pat. No. 7,921,046, the interest list is stored in the form of a bit vector where the position of each bit corresponds to a downstream consumer. For the IFS, a downstream consumer may be a physical output port, a multicast group, a specific host or server, a specific application (such as a feed handler), etc. The scope of a “consumer” depends on the downstream platform architecture. Associated with each consumer is a message queue that contains the messages destined for the consumer. A fair scheduler ensures that each of the message queues receives fair service. Packetization logic reads multiple updates from the selected message queue and packages the updates into a packet for transmission on the prescribed output port, using the prescribed network address and transport port. Messages can be combined into an outgoing Ethernet frame with appropriate MAC-level, and optionally IP-level headers.

(80) Preferably, the packetization logic constructs maximally sized packets: the logic reads as many messages as possible from the queue until the maximum packet size is reached or the message queue is empty. Note that packetization strategy and destination parameters may be specified via packaging parameters stored in a table. The packetization logic simply performs a lookup using the queue number that it is currently servicing in order to retrieve the appropriate parameters. The interest list and packaging parameter tables are preferably managed by configuration, control, and table management logic hosted on a co-resident control processor.

(81) Note that the messages in the newly constructed packets may have been transmitted by their concomitant feed sources in different packets or in the same packet with other messages that are now excluded. This is an example of the IFS constructing a customized “feed” for downstream consumers.

(82) If downstream consumers are equipped with network interface devices that allow for custom protocol implementation, e.g. an FPGA connected directly to the physical network link, then additional optimizations may be implemented by the packetization logic. For example, the Ethernet MAC-level (and above) headers and CRC trailer may be stripped off any packet. By doing so, unnecessary overhead can be removed from packets, reducing packet sizes, reducing data transmission latency, and reducing the amount of processing required to consume the packets. As shown in FIG. 16, this optimization may apply to latency-sensitive trading applications, feed handlers, or ticker plants.

(83) 7. Depth Price Aggregation and Synthetic Quotes:

(84) With sufficient processing and memory resources, additional data normalization functions may be performed by the IFS, and thus offloaded from platform components such as feed handlers, virtual order book engines, and ticker plants. One such function is price-normalization for order-based depth of market feeds. As described in the above-referenced and incorporated U.S. Pat. No. 7,921,046, WO Pub. WO 2010/077829, and U.S. patent application Ser. No. 13/316,332, a number of market data feeds operate at the granularity of individual orders to buy or sell a financial instrument. The majority of real-time updates represent new orders, modifications to existing orders, or deletions of existing orders. As described in these incorporated references, a significant number of market data applications choose to consume the order-based depth of market feeds simply due to the reduced data delivery latency relative to top-of-book or consolidated feeds. However, the applications typically do not require visibility into the individual orders, but rather choose to view pricing information as a limited-depth, price-aggregated book, or as a top-of-book quote. In the above-referenced and incorporated U.S. Pat. No. 7,921,046, WO Pub. WO 2010/077829, and U.S. patent application Ser. No. 13/316,332, a number of techniques are disclosed for efficiently performing price aggregation in parallel processing elements such as reconfigurable hardware devices. The same methods can be applied in the context of an intelligent feed switch to offload price aggregation from downstream consumers. For example, rather than consuming the NASDAQ Totalview feed in its raw order-referenced format, downstream consumers can consume price-aggregated updates reflecting new price points, changes to existing price points, and deletions of price points from the book. This can reduce the number of update events to downstream consumers.

(85) Note that price aggregation may be performed on a per-symbol, per-market basis (e.g. NASDAQ market only), or on a per-symbol, pan-market basis (e.g. NASDAQ, NYSE, BATS, ARCA, Direct Edge) to facilitate virtual order book views.

(86) A further reduction in the number of updates consumed by downstream consumers can be achieved by performing size filtering. Size filtering is defined as the suppression of an update if the result of the update is a change in aggregate volume (size) at a pre-existing price point, where the amount of the change relative to the most recent update transmitted to consumers is less than a configured threshold. Note that the threshold may be relative to the current volume, e.g. a change in size of 50%.

(87) Again, if sufficient processing and memory resources are deployed within the IFS, a synthetic quote engine can be included. As described in the above-referenced and incorporated U.S. Pat. No. 7,921,046, WO Pub. WO 2010/077829, and U.S. patent application Ser. No. 13/316,332, price-aggregated entries can be sorted into a price book view for each symbol. The top N levels of the price-aggregated represent a top-of-book quote. Note that N is typically one (i.e. only the best bid and offer values), but N may be set to be a small value such as three (3) to enhance the quote with visibility into the next N-1 price levels in the book. The techniques described in these incorporated references can be used to efficiently sort price-aggregated updates into price books and generate top-of-book quotes when an entry in the top N

levels changes using parallel processing resources.

(88) 8. Event Caching:

(89) As previously described, the IFS is capable of only transmitting updates for symbols for which downstream consumers are interested using the symbol-based routing described above. If a consumer wishes to add a symbol to its set of interest, the consumer would need to wait until a subsequent quote event is transmitted by the feed source in order to receive the current pricing for the associated financial instrument. A simple form of a cache can be efficiently implemented in the IFS in order to allow downstream consumers to immediately receive current pricing data for a financial instrument if its symbol is dynamically added to its set of interest during a trading session. For feeds that provide top-of-book quote updates and last trade reports, the IFS can maintain a simply last event cache that stores the most recent quote and most recent trade event received on a per-symbol, per-market basis. Specifically, a table of events is maintained where an entry is located using the symbol index, MIC, and MSIC. When the set of interest changes for a given downstream consumer, the current quote and trade events in the event cache are transmitted to the consumer. This allows the consumer to receive the current bid, offer, and last traded price information for the instrument.

(90) If sufficient processing resources exist in the IFS, a full last value cache (LVC) can be maintained as described in the above-referenced and incorporated U.S. Pat. No. 7,921,046.

(91) 9. Data Quality Monitoring:

(92) The IFS can be also be configured to monitor a wide variety of data quality metrics on a per-symbol, per-market basis. A list of data quality metrics includes but is not limited to: Line gap: packet loss experienced on the line carrying updates for the symbol. Line dead: the input feed line is detected to be in a “dead” state where no data is being received. Locked market: the best bid and offer prices for the instrument on the given market are identical. Crossed market: the best bid price is larger than the best offer price for the instrument on the given market

(93) The data quality can be reflected in an enumerated value and included in messages transmitted to downstream consumers as an appended field, as previously described. These enumerated data quality states can be used by the IFS and/or downstream consumers to perform a variety data quality mitigation operations.

(94) 10. Data Source Failover:

(95) An example of a data quality mitigation operation is to provide data source failover. As previously described, there may be multiple data sources for market data updates from a given market, hence the need for a data source identification code (DSIC). Rather specify a specific <symbol, market, data source> tuple when establishing interest in an instrument, downstream consumers may specify a <symbol, market> tuple where the “best” data source is selected by the IFS. A prioritized list of data sources for each market is specified in the control logic. When the data quality associated with the current preferred data source for a market transitions to “poor” quality state, the IFS automatically transitions to the next highest-priority data source for the market. The data quality states that constitute “poor” quality are configured in the control logic. When a data source transition occurs, the control logic alters the interest list entries associated with affected instruments and downstream consumers. Note that if a higher-priority data source transitions out of a “poor” quality state, the IFS automatically transitions back to the higher-priority data source. Preferably, the IFS is configured to apply hysteresis to the data source failover function to prevent thrashing between data sources. Note that data source failover may rely on the presence of other functions within the IFS such as synthetic quote generation if failover is to be supported between depth of market feeds and top-of-book quote feeds.

(96) 11. Monitoring, Configuration, and Control:

(97) The monitoring, configuration, and control logic described is preferably hosted on a co-resident processor in the IFS. This logic may interface with applications in the electronic platform or remote operations applications. In one embodiment of the IFS, control messages are received from an egress port. This allows one or more applications in the electronic trading platform to specify symbol routing parameters, packet and message parsing templates, prioritized lists of data sources, gap reporting and mitigation parameters, etc.

(98) In addition, a variety of statistics counters and informational registers are maintained by the offload engines that can be accessed by the control logic in the IFS such as per-line packet and message

counters, packet and message rates, gap counters and missing sequence registers, packet size statistics, etc. These statistics are made available to the external world via common mechanisms in the art, including SNMP, HTML, etc.

(99) 12. Feed Generation:

(100) The IFS can also be used by feed sources (exchanges and consolidated feed vendors) to offload many of the functions required in feed generation. These tasks are largely the inverse of those performed by feed consumers. Specifically, the IFS can be configured to encode updates using prescribed encoding templates and transmit the updates on specified multicast groups, output ports, etc. Other functions that are applicable to feed generation include high-resolution timestamping, rate monitoring, and data quality monitoring.

(101) C. Intelligent Distribution Switch:

(102) The same methods and apparatuses can be applied to the task of distributing data throughout the electronic trading platform. As shown in FIG. 24, an Intelligent Distribution Switch (IDS) can be positioned downstream of market data normalization components in the electronic trading platform. The IDS can be used to offload distribution functions from normalization components such as ticker plants, to offload data consumption and management functions from downstream consumers such as trading applications, and to introduce new capabilities into the distribution network in the electronic trading platform. Examples of distribution capabilities are described in the above-referenced and incorporated U.S. Pat. App. Ser. No. 61/570,670.

(103) The IDS architecture can be one of the previously described variants shown in FIGS. 21, 22, and 23. Note that the number of switch ports and amount of interconnect bandwidth between internal devices (FPGAs, switch ASICs, memory, etc.) may be provisioned differently for an IDS application, relative to an IFS application.

(104) As shown in FIG. 25, the IDS may host one or more distribution functions. The IDS can be used to offload the task of interest-based distribution. The IDS can maintain a mapping from instrument symbol to interest list, an example of such a mapping being described in the above-referenced and incorporated U.S. Pat. No. 7,921,046. If point-to-point transmission protocols are in use, then the IDS makes the requisite copies of the update event and addresses each event for the specified consumer. By offloading this function, upstream components such as ticker plants only need to propagate a single copy of each update event. This reduces the processing resource requirement, or allows the processing resources previously dedicated to interest list maintenance and event replication to be redeployed for other purposes.

(105) Data source failover may also be performed by the IDS. Like the previously described data source failover function performed in the IFS, the IDS allows downstream consumers to specify a prioritized list of normalized data sources. When the preferred source becomes unavailable or the data quality transitions to an unacceptable state, the IDS switches to the next highest priority normalized data source.

(106) The IDS may also perform customized computations on a per-consumer basis. Example computations include constructing user-defined Virtual Order Books, computing basket computations, computing options prices (and implied volatilities) and generating user-defined Best Bid and Offer (BBO) quotes (see the above-referenced and incorporated U.S. Pat. Nos. 7,840,482 and 7,921,046, U.S. Pat. App. Pub. 2009/0182683, and WO Pub. WO 2010/077829 for examples of hardware-accelerated processing modules for such tasks). By performing these functions in an IDS at the “edge” of the distribution network allows the functions to be customized on a per consumer basis. Note that a ticker plant distributing data to hundreds of consumers may not have the processing capacity to perform hundreds of customized computations, one for each consumer. Examples of other customized per consumer computations include: liquidity target Net Asset Value (NAV) computations, future/spot price transformations, and currency conversions.

(107) Additionally, the IDS may host one or more of the low latency data distribution functions described in the above-referenced and incorporated U.S. Pat. App. Ser. No. 61/570,670. In one embodiment, the IDS may perform all of the functions of an Edge Cache. In another embodiment, the IDS may perform all of the functions of a Connection Multiplexer. As such, the IDS includes at least one instance of a multi-class distribution engine (MDE) that includes some permutation of Critical Transmission Engine, Adaptive Transmission Engine, or Metered Transmission Engine.

(108) Like the customized per consumer computations, the IDS may also perform per consumer protocol bridging. For example, the upstream connection from the IDS to a ticker plant may use a point-to-point Remote Direct Memory Access (RDMA) protocol. The IDS may be distributing data to a set of consumers via point-to-point connections using the Transmission Control Protocol (TCP) over Internet Protocol (IP), and distributing data to another set of consumers via a proprietary reliable multicast protocol over Unreliable Datagram Protocol (UDP).

(109) 1. Low Overhead Communication Protocols:

(110) Note that if intelligent FPGA NICs are used in the consuming machines, then a direct FPGA-to-FPGA wire path exists between FPGA in the Switch and the FPGA in the NIC. This eliminates the need for Ethernet frame headers, IP headers, CRCs, inter-frame spacing and other overhead, and allows the FPGA in the switch to communicate directly with the FPGA in the NIC, without being constrained to specific communication protocols.

(111) While the present invention has been described above in relation to exemplary embodiments, various modifications may be made thereto that still fall within the invention's scope, as would be recognized by those of ordinary skill in the art. Such modifications to the invention will be recognizable upon review of the teachings herein. As such, the full scope of the present invention is to be defined solely by the appended claims and their legal equivalents.

Claims

1. A system for processing a plurality of data packets from a plurality of financial market data feeds, wherein each of a plurality of the data packets comprises (1) a destination address that identifies a destination for the data packet, (2) a source address that identifies a source for the data packet, (3) a protocol identifier that identifies a protocol for the data packet, and (4) a packet payload, the packet payload comprising plurality of financial market data messages, the financial market data messages comprising a plurality of fields of financial market data pertaining to a plurality of financial instruments, the system comprising: a network element of a data distribution network through which data packets are communicated to an electronic trading platform, wherein the network element comprises a packet switch, a repeater, a router, or a firewall, the network element comprising a processor, the processor configured to: receive the data packets; and for each of a plurality of the received data packets, (1) map the received data packet to parsing metadata applicable to the received data packet based on the received data packet's destination address, source address, and protocol identifier and (2) append a field to the received data packet, wherein the appended field comprises the parsing metadata; and wherein the network element is configured to communicate the data packets to the destinations corresponding to the destination addresses of the data packets, wherein the communicated data packets include the appended fields, the network element thereby being configured to offload tasks for mapping parsing metadata to data packets from the electronic trading platform.
2. The system of claim 1 wherein the parsing metadata includes an identifier for the financial market data feed for the data packet to which the parsing metadata is applicable.
3. The system of claim 2 wherein the parsing metadata further includes an identifier for a line within the financial market data feed for the data packet to which the parsing metadata is applicable.
4. The system of claim 1 wherein the parsing metadata includes an identifier for an exchange that pertains to the data packet to which the parsing metadata is applicable.
5. The system of claim 1 wherein the parsing metadata includes a packet parsing template that identifies how the data packet is to be parsed.
6. The system of claim 1 wherein the parsing metadata includes a message parsing template that identifies how the financial market data messages of the data packet are to be parsed into their constituent fields of financial market data.
7. The system of claim 1 wherein the processor is further configured to, for each of a plurality of the received data packets, map the received data packet to its corresponding parsing template using a hash key and a hash table, wherein the hash key is based on the destination address, the source address, and the protocol identifier of the received data packet.

8. The system of claim 1 wherein the processor comprises a field programmable gate array (FPGA).
 9. The system of claim 1 wherein the processor comprises an application-specific integrated circuit (ASIC).
 10. The system of claim 1 wherein the processor comprises a graphics processing unit (GPU).
 11. The system of claim 1 wherein the processor comprises a chip multi-processor (CMP).
 12. The system of claim 1 wherein the network element comprises the packet switch, wherein the packet switch comprises (1) a plurality of ports for receiving and communicating the data packets, (2) a switching fabric, and (3) the processor.
 13. The system of claim 12 wherein the packet switch provides 10 Gigabit Ethernet connectivity for receiving and communicating the data packets via a plurality of the ports.
 14. The system of claim 1 further comprising: a ticker plant configured to receive the communicated data packets and parse the received communicated data packets into a plurality of parsed fields based on the parsing metadata included in the appended fields, wherein the ticker plant comprises a field programmable gate array (FPGA), the FPGA configured to generate composite views of a plurality of financial instruments based on a plurality of the parsed fields.
 15. The system of claim 1 further comprising: a ticker plant configured to receive the communicated data packets and parse the received communicated data packets into a plurality of parsed fields based on the parsing metadata included in the appended fields, wherein the ticker plant comprises a field programmable gate array (FPGA), the FPGA configured to perform last value cache (LVC) updates for a plurality of financial instruments based on a plurality of the parsed fields.
 16. A method for processing a plurality of data packets from a plurality of financial market data feeds, the method comprising: receiving a data packet by a processor within a network element of a data distribution network through which data packets are communicated to an electronic trading platform, wherein the network element comprises a packet switch, a repeater, a router, or a firewall, wherein the received data packet comprises (1) a destination address that identifies a destination for the data packet, (2) a source address that identifies a source for the data packet, (3) a protocol identifier that identifies a protocol for the data packet, and (4) a packet payload, the packet payload comprising plurality of financial market data messages, the financial market data messages comprising a plurality of fields of financial market data pertaining to a plurality of financial instruments; the processor mapping the data packet to parsing metadata applicable to the data packet based on the data packet's destination address, source address, and protocol identifier; the processor appending a field to the data packet, wherein the appended field comprises the parsing metadata; and the network element communicating the data packet to the destination corresponding to the destination address of the data packet, wherein the communicated data packet includes the appended field, the network element thereby offloading tasks for mapping parsing metadata to data packets from the electronic trading platform; and performing the (1) receiving step, (2) the mapping step, (3) the appending step, and (4) and communicating step for a plurality of data packets from a plurality of financial market data feeds.
 17. The system of claim 1 wherein the network element comprises the repeater.
 18. The system of claim 1 wherein the network element comprises the router.
 19. The system of claim 1 wherein the network element comprises the firewall.
-