

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12417549
B2
September 16, 2025
Kim; Dae Jin et al.

Vanishing point determination, symmetry-based boundary refinement, and component detection for vehicle object detection or other applications

Abstract

A method includes obtaining, using at least one processing device, a vanishing point and a boundary based on image data associated with a scene, where the boundary is associated with a detected object within the scene. The method also includes repeatedly, during multiple iterations and using the at least one processing device, (i) identifying multiple patches within the boundary and (ii) determining a similarity of the image data contained within the multiple patches. The method further includes identifying, using the at least one processing device, a modification to be applied to the boundary based on the identified patches and the determined similarities. In addition, the method includes generating, using the at least one processing device, a refined boundary based on the modification, where the refined boundary identifies a specified portion of the detected object.

Inventors: Kim; Dae Jin (Rancho Palos Verdes, CA), Arft; David R. (Torrance, CA), Choi; Jongmoo (Gardena, CA)

Applicant: Canoo Technologies Inc. (Torrance, CA)

Family ID: 1000008818245

Assignee: CANOO TECHNOLOGIES INC. (Torrance, CA)

Appl. No.: 17/819584

Filed: August 12, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230100507 A1	Mar. 30, 2023

Related U.S. Application Data

us-provisional-application US 63249732 20210929

Publication Classification

Int. Cl.: G06T7/536 (20170101); B60W10/18 (20120101); B60W10/20 (20060101); B60W50/00 (20060101); B60W50/14 (20200101); B60W50/16 (20200101); B60W60/00 (20200101); G06T7/68 (20170101); G06T7/73 (20170101); G06V20/56 (20220101); G06V20/58 (20220101); G06V20/62 (20220101)

U.S. Cl.:

CPC **G06T7/536** (20170101); **B60W10/18** (20130101); **B60W10/20** (20130101); **B60W50/14** (20130101); **B60W50/16** (20130101); **G06T7/68** (20170101); **G06T7/73** (20170101); **G06V20/56** (20220101); **G06V20/58** (20220101); **G06V20/625** (20220101); B60W2050/0004 (20130101); B60W2050/143 (20130101); B60W2050/146 (20130101); B60W60/001 (20200201); B60W2420/403 (20130101); B60W2552/50 (20200201); B60W2554/20 (20200201); B60W2710/18 (20130101); B60W2710/20 (20130101); B60W2720/10 (20130101); B60W2720/106 (20130101); G06T2207/30252 (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10549694	12/2019	Murad	N/A	B60R 1/26
2006/0212215	12/2005	Koulinitch	701/96	G06T 7/564
2010/0097455	12/2009	Zhang	701/25	G06V 20/588
2013/0266175	12/2012	Zhang et al.	N/A	N/A
2014/0314279	12/2013	Zhang et al.	N/A	N/A
2015/0063709	12/2014	Wang	382/199	G06T 7/12
2016/0014406	12/2015	Takahashi	348/148	G01S 5/16
2019/0026905	12/2018	Siegemund	N/A	G06T 7/277
2019/0197693	12/2018	Zagaynov et al.	N/A	N/A
2020/0226805	12/2019	Szeliski et al.	N/A	N/A
2020/0334476	12/2019	Wang	N/A	G06F 18/214
2021/0142497	12/2020	Pugh et al.	N/A	N/A
2021/0217203	12/2020	Kim	N/A	G06T 9/001
2021/0406679	12/2020	Wen	N/A	G06V 20/56

OTHER PUBLICATIONS

International Search Report and Written Opinion of the International Searching Authority dated Jan. 18, 2023, in connection with International Application No. PCT/US2022/074939, 31 pages. cited by applicant

Non-final Office Action dated Jun. 4, 2024, in connection with U.S. Appl. No. 17/819,583, 17 pages. cited by applicant

Kong et al., “Vanishing point detection for road detection”, IEEE Conference on Computer Vision and Pattern Recognition, Jun. 2009, 9 pages. cited by applicant

Ozgunalp et al., “Multiple Lane Detection Algorithm Based on Novel Dense Vanishing Point Estimation”, IEEE Transactions on Intelligent Transportation, vol. 18, No. 3, Mar. 2017, 12 pages. cited by applicant

Kim et al., “Vanishing Point Determination, Symmetry-Based Boundary Refinement, and Component Detection for Vehicle Object Detection or Other Applications”, U.S. Appl. No. 17/819,583, filed Aug. 12, 2022, 50 pages. cited by applicant

Kim et al., “Vanishing Point Determination, Symmetry-Based Boundary Refinement, and Component Detection for Vehicle Object Detection or Other Applications”, U.S. Appl. No. 17/819,587, filed Aug. 12, 2022, 51 pages. cited by applicant

Primary Examiner: Gilliard; Delomia L

Attorney, Agent or Firm: ALLEN, DYER, DOPPELT, GILCHRIST, P.A.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION AND PRIORITY CLAIM (1) This application claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 63/249,732 filed on Sep. 29, 2021. This provisional application is hereby incorporated by reference in its entirety.

TECHNICAL FIELD

(1) This disclosure relates generally to object detection systems. More specifically, this disclosure relates to vanishing point determination, symmetry-based boundary refinement, and component detection for vehicle object detection or other applications.

BACKGROUND

(2) Identifying nearby, moving, or other objects in a scene is often an important or useful function in many autonomous applications, such as in vehicles supporting advanced driving assist system (ADAS) or autonomous driving (AD) features, or other applications. Performing accurate object detection often involves the use of complex computer vision algorithms or complex deep neural networks. Unfortunately, these algorithms tend to require larger amounts of computing resources and can generate results more slowly than desired, particularly for automotive applications.

SUMMARY

(3) This disclosure provides vanishing point determination, symmetry-based boundary refinement, and component detection for vehicle object detection or other applications.

(4) In a first embodiment, a method includes obtaining, using at least one processing device, image data associated with a scene. The method also includes identifying, using the at least one processing device, multiple line segments based on the image data. The method further includes identifying, using the at least one processing device, one or more boundaries around one or more objects detected in the image data. In addition, the method includes estimating, using the at least one processing device, a position of a vanishing point associated with the image data based on multiple collections of the line segments while excluding, from the multiple collections, one or more of the line segments that overlap with or that are included within the one or more boundaries.

(5) In a second embodiment, an apparatus includes at least one processing device configured to obtain image data associated with a scene. The at least one processing device is also configured to identify multiple line segments based on the image data. The at least one processing device is further configured to identify one or more boundaries around one or more objects detected in the image data. In addition, the at least one processing device is configured to estimate a position of a vanishing point associated with the image data based on multiple collections of the line segments while excluding, from the multiple collections, one or more of the line segments that overlap with or that are included within the one or more boundaries.

(6) In a third embodiment, a non-transitory machine-readable medium contains instructions that when executed cause at least one processor to obtain image data associated with a scene. The medium also contains instructions that when executed cause the at least one processor to identify multiple line segments based on the image data. The medium further contains instructions that when executed cause the at least one processor to identify one or more boundaries around one or more objects detected in the image data. In addition, the medium contains instructions that when executed cause the at least one processor to estimate a position of a vanishing point associated with the image data based on multiple collections of the line segments while excluding, from the multiple collections, one or more of the line segments that overlap with or that are included within the one or more boundaries.

(7) In a fourth embodiment, a method includes obtaining, using at least one processing device, a vanishing point and a boundary based on image data associated with a scene, where the boundary is associated with a detected object within the scene. The method also includes repeatedly, during multiple iterations and using the at least one processing device, (i) identifying multiple patches within the boundary and (ii) determining a similarity of the image data contained within the multiple patches. The method further includes identifying, using the at least one processing device, a modification to be applied to the boundary based on the identified patches and the determined similarities. In addition, the method includes generating, using the at least one processing device, a refined boundary based on the modification, where the refined boundary identifies a specified portion of the detected object.

(8) In a fifth embodiment, an apparatus includes at least one processing device configured to obtain a vanishing point and a boundary based on image data associated with a scene, where the boundary is associated with a detected object within the scene. The at least one processing device is also configured to repeatedly, during multiple iterations, (i) identify multiple patches within the boundary and (ii) determine a similarity of the image data contained within the multiple patches. The at least one processing device is further configured to identify a modification to be applied to the boundary based on the identified patches and the determined similarities. In addition, the at least one processing device is configured to generate a refined boundary based on the modification, where the refined boundary identifies a specified portion of the detected object.

(9) In a sixth embodiment, a non-transitory machine-readable medium contains instructions that when executed cause at least one processor to obtain a vanishing point and a boundary based on image data associated with a scene, where the boundary is associated with a detected object within the scene. The medium also contains instructions that when executed cause the at least one processor to repeatedly, during multiple iterations, (i)

identify multiple patches within the boundary and (ii) determine a similarity of the image data contained within the multiple patches. The medium further contains instructions that when executed cause the at least one processor to identify a modification to be applied to the boundary based on the identified patches and the determined similarities. In addition, the medium contains instructions that when executed cause the at least one processor to generate a refined boundary based on the modification, where the refined boundary identifies a specified portion of the detected object.

(10) In a seventh embodiment, a method includes obtaining, using at least one processing device, a refined boundary identifying a specified portion of a detected object within a scene, where the refined boundary is associated with image data. The method also includes repeatedly, during multiple iterations and using the at least one processing device, (i) identifying multiple regions within the refined boundary and (ii) determining a similarity of the image data contained within the multiple regions. In addition, the method includes identifying, using the at least one processing device, one or more locations of one or more components of the detected object based on the identified regions and the determined similarities.

(11) In an eighth embodiment, an apparatus includes at least one processing device configured to obtain a refined boundary identifying a specified portion of a detected object within a scene, where the refined boundary is associated with image data. The at least one processing device is also configured to repeatedly, during multiple iterations, (i) identify multiple regions within the refined boundary and (ii) determine a similarity of the image data contained within the multiple regions. The at least one processing device is further configured to identify one or more locations of one or more components of the detected object based on the identified regions and the determined similarities.

(12) In a ninth embodiment, a non-transitory machine-readable medium contains instructions that when executed cause at least one processor to obtain a refined boundary identifying a specified portion of a detected object within a scene, where the refined boundary is associated with image data. The medium also contains instructions that when executed cause the at least one processor to repeatedly, during multiple iterations, (i) identify multiple regions within the refined boundary and (ii) determine a similarity of the image data contained within the multiple regions. The medium further contains instructions that when executed cause the at least one processor to identify one or more locations of one or more components of the detected object based on the identified regions and the determined similarities.

(13) Other technical features may be readily apparent to one skilled in the art from the following figures, descriptions, and claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) For a more complete understanding of this disclosure and its advantages, reference is now made to the following description taken in conjunction with the accompanying drawings, in which like reference numerals represent like parts:

(2) FIG. 1 illustrates an example system supporting vanishing point determination according to this disclosure;

(3) FIG. 2 illustrates an example method for vanishing point determination according to this disclosure;

(4) FIG. 3 illustrates an example vanishing point determination according to this disclosure;

(5) FIG. 4 illustrates an example system supporting symmetry-based boundary refinement according to this disclosure;

(6) FIG. 5 illustrates an example method for symmetry-based boundary refinement according to this disclosure;

(7) FIG. 6 illustrates an example symmetry-based boundary refinement according to this disclosure;

(8) FIG. 7 illustrates an example system supporting component detection according to this disclosure;

(9) FIG. 8 illustrates an example method for component detection according to this disclosure;

(10) FIG. 9 illustrates an example component detection according to this disclosure;

(11) FIG. 10 illustrates an example usage of vanishing point determination, symmetry-based boundary refinement, and/or component detection according to this disclosure;

(12) FIG. 11 illustrates an example design flow for employing one or more tools to design hardware that implements one or more control functions according to this disclosure; and

(13) FIG. 12 illustrates an example device supporting execution of one or more tools to design hardware that implements one or more control functions according to this disclosure.

DETAILED DESCRIPTION

(14) FIGS. 1 through 12, described below, and the various embodiments used to describe the principles of this disclosure are by way of illustration only and should not be construed in any way to limit the scope of this disclosure. Those skilled in the art will understand that the principles of this disclosure may be implemented in any

type of suitably arranged device or system.

(15) As noted above, identifying nearby, moving, or other objects in a scene is often an important or useful function in many autonomous applications, such as in vehicles supporting advanced driving assist system (ADAS) or autonomous driving (AD) features, or other applications. Performing accurate object detection often involves the use of complex computer vision algorithms or complex deep neural networks. Unfortunately, these algorithms tend to require larger amounts of computing resources and can generate results more slowly than desired, particularly for automotive applications.

(16) In one aspect, this disclosure provides techniques for determining vanishing points in images of scenes associated with a vehicle (often referred to as a target vehicle). As described in more detail below, an image of a scene associated with a target vehicle (such as an image of the scene in front of the target vehicle) can be captured and processed to identify the vanishing point in the image. In an image containing a set of parallel or substantially parallel lines in three-dimensional space (such as lane-marking lines), the lines would form a generally triangular shape in the image plane due to the perspective of a device capturing the image. A vanishing point is defined as the point at which the lines in the set converge within the image plane. The determination of the vanishing point in an image is useful in various image processing applications, such as in algorithms used for line or lane estimation, indirect measurements of vehicle pitching angles, and identification of ego/non-ego lane driving vehicles.

(17) In another aspect, this disclosure provides techniques for performing symmetry-based boundary refinement in order to refine one or more boundaries for one or more objects detected in a scene around a target vehicle. As described in more detail below, once a bounding box or other boundary roughly identifying another vehicle is determined in an image of a scene around a target vehicle, a symmetry-based approach can be used to locate a more-accurate boundary around a portion of the other vehicle (such as around a rear of the other vehicle) based on an expected symmetry of the other vehicle. For example, another vehicle driving ahead of or towards the target vehicle is expected to have a generally-symmetrical shape, which can be used to identify a more-accurate boundary around a portion of the other vehicle in the image. In some cases, to avoid reliance on the other vehicle's specific make/type/model, images may be subjected to no or minimal enhancements (such as only Sobel edge detection or other edge detection) without any pre-trained patch/template matching. The determination of more-accurate vehicle boundaries can be useful in various applications, such as identifying a direction of travel or a change in the direction of travel of the other vehicle relative to the target vehicle or identifying a surface of the other vehicle to be used for depth estimation.

(18) In still another aspect, this disclosure provides techniques for performing component detection in order to identify one or more specific components on the rear portion(s) or other portion(s) of at least one other vehicle in a scene around a target vehicle. As described in more detail below, once the boundary of another vehicle is determined and refined, the image data associated with the other vehicle within the refined boundary can be analyzed to identify one or more regions of the other vehicle associated with one or more specific components of the other vehicle, such as rear taillights or a license plate of the other vehicle. Again, this can be accomplished without prior knowledge of other vehicle's specific make/type/model. This type of information can be useful in various applications, such as identifying a direction of travel or a change in the direction of travel of the other vehicle relative to the target vehicle or identifying a surface of the other vehicle to be used for depth estimation.

(19) In this way, these techniques can be used to identify various information (such as a vanishing point, a symmetry-based boundary, and/or detected vehicle components) that is useful in performing one or more functions, such as one or more ADAS or AD functions. Moreover, this information can be determined using less computing resources and in faster times, which enables use of this information in resource-constrained applications and in real-time applications.

(20) FIG. 1 illustrates an example system **100** supporting vanishing point determination according to this disclosure. In this particular example, the system **100** takes the form of an automotive vehicle, such as an electric vehicle. However, any other suitable system may support vanishing point determination, such as other types of vehicles, autonomous robots, or other autonomous or non-autonomous systems.

(21) As shown in FIG. 1, the system **100** includes at least one processor **102** configured to control one or more operations of the system **100**. In this example, the processor **102** may interact with one or more sensors **104** and with one or more components coupled to a bus **106**. In this particular example, the one or more sensors **104** include one or more cameras **104a**, or other imaging sensors, and the bus **106** represents a controller area network (CAN) bus. However, the processor **102** may interact with any additional sensor(s) and communicate over any other or additional bus(es).

(22) The one or more cameras **104a** are configured to generate images of scenes around the system **100**. Note that other or additional types of sensors may be used here, such as one or more radio detection and ranging (RADAR) sensors, light detection and ranging (LIDAR) sensors, other types of imaging sensors, or inertial measurement units (IMUs). Measurements or other data from the sensors **104** are used by the processor **102** or other component(s) as described below to perform various functions. In some cases, the sensors **104** may include a

single camera **104a**, such as one camera positioned on the front of a vehicle. In other cases, the sensors **104** may include multiple cameras **104a**, such as one camera positioned on the front of a vehicle, one camera positioned on the rear of the vehicle, and two cameras positioned on opposite sides of the vehicle.

(23) The processor **102** can process the information from the sensors **104** in order to detect objects around or proximate to the system **100**, such as one or more vehicles, obstacles, or people near the system **100**. The processor **102** can also process the information from the sensors **104** in order to perceive lane-marking lines or other markings on a road, floor, or other surface. The processor **102** can further use various information to generate predictions associated with the system **100**, such as to predict the future path(s) of the system **100** or other vehicles, identify a center of a lane in which the system **100** is traveling, or predict the future locations of objects around the system **100**.

(24) In this example, the processor **102** performs an object detection function **108**, which generally operates to identify objects around the system **100** in a real-time manner based on images or other measurements from the sensor(s) **104**. For example, the object detection function **108** can use images from one or more cameras **104a** or data from other sensors **104** to identify external objects around the system **100**, such as other vehicles moving around or towards the system **100** or pedestrians or objects near the system **100**. The object detection function **108** can also identify one or more characteristics of each of one or more detected objects, such as an object class (a type of object) and a boundary around the detected object. The processor **102** may use any suitable technique to identify the objects around the system **100** based on data from the sensor(s) **104**. Various techniques for object detection are known in the art, and additional techniques for object detection are sure to be developed in the future. Any of these techniques may be used by the processor **102** here to implement the object detection function **108**.

(25) The processor **102** also performs a line segment detection function **110**, which generally operates to identify line segments within the images from one or more cameras **104a** or the data from other sensors **104**. Each detected line segment may generally represent a portion in an image or other sensor data identifying a relatively straight feature. As particular examples, detected line segments may represent portions of lane-marking lines or other markings on a road, floor, or other surface; edges of nearby buildings or other nearby structures; or edges or other features of other objects captured in images or other data. The processor **102** may use any suitable technique to identify line segments based on data from the sensor(s) **104**. Various techniques for line segment detection are known in the art, and additional techniques for line segment detection are sure to be developed in the future. Any of these techniques may be used by the processor **102** here to implement the line segment detection function **110**.

(26) The processor **102** further performs an incremental aggregation function **112** and a vanishing point detection function **114**. The incremental aggregation function **112** generally operates to identify different collections of the line segments detected by the line segment detection function **110** and to estimate a crossing point of the line segments for each collection (which may also involve determination of a residual value). The vanishing point detection function **114** generally operates to calculate a vanishing point for the image based on the determined crossing points, such as by using the residual values as weights during a weighted combination of the crossing points. Example operations that can be performed by the incremental aggregation function **112** and the vanishing point detection function **114** are described below. The final vanishing point determined for a scene may be used in any suitable manner, including in the techniques described below.

(27) Note that the functions **108-114** shown in FIG. 1 and described above may be implemented in any suitable manner in the system **100**. For example, in some embodiments, various functions **108-114** may be implemented or supported using one or more software applications or other software instructions that are executed by at least one processor **102**. In other embodiments, at least some of the functions **108-114** can be implemented or supported using dedicated hardware components. In general, the functions **108-114** described above may be performed using any suitable hardware or any suitable combination of hardware and software/firmware instructions.

(28) The processor **102** itself may also be implemented in any suitable manner, and the system **100** may include any suitable number(s) and type(s) of processors or other processing devices in any suitable arrangement. Example types of processors **102** that may be used here include one or more microprocessors, microcontrollers, digital signal processors (DSPs), application specific integrated circuits (ASICs), field programmable gate arrays (FPGAs), or discrete circuitry. Each processor **102** may also have any suitable number of processing cores or engines. In some cases, multiple processors **102** or multiple processing cores or engines in one or more processors **102** may be used to perform the functions **108-114** described above. This may allow, for instance, the processor(s) **102** to be used to process multiple images and other sensor data in parallel or to perform various operations described in this patent document above and below in parallel.

(29) Although FIG. 1 illustrates one example of a system **100** supporting vanishing point determination, various changes may be made to FIG. 1. For example, various functions and components shown in FIG. 1 may be combined, further subdivided, replicated, omitted, or rearranged and additional functions and components may be added according to particular needs. Also, as noted above, the functionality for vanishing point determination may be used in any other suitable system, and the system may or may not relate to automotive vehicles or other

vehicles.

(30) FIG. 2 illustrates an example method **200** for vanishing point determination according to this disclosure. The method **200** may, for example, be performed using the components of the system **100** shown in FIG. 1. Note, however, that the method **200** may be performed using any other suitable device or system. Also, during the discussion of the method **200**, reference is made to FIG. 3, which illustrates an example vanishing point determination according to this disclosure. The example of the determination shown in FIG. 3 is for illustration only and is merely meant, to illustrate how various steps in FIG. 2 may be performed.

(31) As shown in FIG. 2, an image of a scene around a vehicle is obtained at step **202**. This may include, for example, the processor **102** obtaining at least one image from at least one camera **104a** or obtaining other sensor data from at least one other type of sensor **104**. One or more objects in the scene are identified using the image at step **204**. This may include, for example, the processor **102** performing the object detection function **108** to identify at least one object and a bounding box or other boundary around each object. Line segments in the scene are identified using the image at step **206**. This may include, for example, the processor **102** performing the line segment detection function **110** to identify segments of lane-marking lines, other manmade lines, edges of buildings or other structures, or other line segments in the image.

(32) An example of this is shown in FIG. 3, where an image **300** is processed to identify two bounding boxes **302**, **304** associated with two objects and to identify three line segments **306**, **308**, **310**. Bounding boxes may be referred to below using the notation $\{\text{right arrow over } (\theta)\}.\text{sub.bbox}, k=y.\text{sub.k}, w.\text{sub.k}, h.\text{sub.k}\}$, where $k=1, \dots, N$. Here, N refers to the total number of bounding boxes, $(x.\text{sub.}\alpha, y.\text{sub.}\alpha)$ refers to the coordinates of the pixel at the top left corner of the $k.\text{sup.th}$ bounding box, and $(w.\text{sub.k}, h.\text{sub.k})$ refers to the width and height of the $k.\text{sup.th}$ bounding box in pixels. Also, line segments may be referred to below using the notation

, where =1, ..., M. Here, M refers to the total number of line segments, and each line segment may be parameterized, such as by using the form =ax+by+c or by using the form =s {right arrow over (p)}.sub.A+(1-s){right arrow over (p)}.sub.B, $s \in [0,1]$.

(33) An iterative process occurs to identify the crossing points for different combinations of the identified line segments. A reference horizon in the image is identified at step **208**. This may include, for example, the processor **102** performing the incremental aggregation function **112** to identify a generally horizontal line around the middle of the captured image where it appears the horizon is located. An example of this is shown in FIG. 3, where a line **312** represents an identified reference horizon.

(34) A set of line segments under the reference horizon is selected at step **210**. This may include, for example, the processor **102** performing the incremental aggregation function **112** to identify a set of the line segments that are present under the identified reference horizon in the image. In some embodiments, the processor **102** can select a set of L line segments that satisfy the following:

$$S.\text{sub.l} = \arg \left\{ \left| \left\{ \left\{ \text{right arrow over } (p) \right\}.\text{sub.A}(y) < h.\text{sub.ref} \text{ and } \left\{ \text{right arrow over } (p) \right\}.\text{sub.B}(y) < h.\text{sub.ref} \right\} \right| = L \right\} \quad (1)$$

Here, $h.\text{sub.ref}$ refers to the identified reference horizon, and $\{\text{right arrow over } (p)\}.\text{sub.A}(y)$ and $\{\text{right arrow over } (p)\}.\text{sub.B}(y)$ refer to the y-coordinates of the ends of the associated line segment. Also, in some embodiments, one or more bounding boxes can be used to filter or ignore line segments that overlap or that are included within the bounding box(es), which can be done to help prevent noisy estimations. This can be expressed as follows:

(35)

$$S_l = \arg \left\{ \left| \left\{ \left\{ \text{right arrow over } (p) \right\} \in \left\{ \text{right arrow over } (\theta) \right\}.\text{sub.bbox} \mid \left\{ \text{right arrow over } (p) \right\}.\text{sub.A} \in \left\{ \text{right arrow over } (\theta) \right\}.\text{sub.bbox}, q \text{ and } \left\{ \text{right arrow over } (p) \right\}.\text{sub.B} \in \left\{ \text{right arrow over } (\theta) \right\}.\text{sub.bbox}, q \forall q = 1, \dots, N \right\} \right| = L \right\} \quad (2)$$

Here, $\{\text{right arrow over } (p)\} \in \{\text{right arrow over } (\theta)\}.\text{sub.bbox}$ represents a predefined operator to check if a point $\{\text{right arrow over } (p)\}$ is contained in a bounding box $\{\text{right arrow over } (\theta)\}.\text{sub.bbox}$.

(36) A crossing point of the line segments in the select set is estimated at step **212**. This may include, for example, the processor **102** performing the incremental aggregation function **112** to identify a crossing point for the L line segments in the set using a “least squares” calculation. In some embodiments, the crossing point may be identified in the following manner:

$$(37) \hat{a}X = (A^T A)^{-1} A^T B \quad (3) \text{ where: } A = \begin{bmatrix} a_1 & b_1 & c_1 \\ \vdots & \vdots & \vdots \\ a_L & b_L & c_L \end{bmatrix}, X = [x \ y \ 1]^T, B = [0 \ 0 \ 0]^T \quad (4)$$

Here, $a.\text{sub.1}-c.\text{sub.1}$ represent a vector of values associated with the first line segment, and $a.\text{sub.L}-c.\text{sub.L}$ represent a vector of values associated with the $L.\text{sup.th}$ line segment. The least squares calculation or other calculation may identify a crossing point $\{\text{right arrow over } (p)\}.\text{sub.cross} = (\hat{x}, \hat{y})$ and a residual r . The crossing point and its associated residual can be stored at step **214**. In some cases, the crossing point and its associated residual can be stored as a pair $S.\text{sub.p} \leftarrow \{\{\text{right arrow over } (p)\}.\text{sub.cross}, r\}$ in a cache or other

memory of the processor **102** or other suitable storage.

(38) A determination is made whether to repeat another iteration at step **216**. This may include, for example, the processor **102** performing the incremental aggregation function **112** to determine whether a user-specified number or other number of estimated crossing points have been identified. If not, the process can return to step **208**. Otherwise, a vanishing point is determined using the stored crossing points and residuals at step **218**. This may include, for example, the processor **102** performing the vanishing point detection function **114** to calculate the vanishing point for the image using the stored crossing points and residuals. In some embodiments, the vanishing point detection function **114** determines the vanishing point using a weighted combination of the crossing points, where the residuals are used as the weights for the combination. This can be expressed as:

$$(39) \quad \overset{\text{fwdarw.}}{p}_{\text{vanish}} = \text{.Math.}_{i=1, \text{.Math.}, K} \tilde{r}_i \overset{\text{fwdarw.}}{p}_{\text{cross}}, \tilde{r}_i = \frac{r_i}{\text{.Math.}_{j=1, \text{.Math.}, K} r_j} \quad (5)$$

Here, K represents the number of crossing points and residuals generated during K iterations of steps **208-214**. As shown in FIG. **3**, a vanishing point **314** may be identified for the image **300**. The identified vanishing point **314** may be used in various ways, including in the techniques described below.

(40) In this way, vanishing point estimation may occur via a progressive aggregation of line segments. This approach can be easily implemented using less-heavy computing resources and/or can be performed in less time. Many of the line segments may be identified using painted lane-marking lines and other painted lines, while other environmental line segments (such as from buildings, road curbs, or urban poles) can also be used. To prevent inconsistent vanishing point estimation, line segments associated with other vehicles may optionally be filtered, such as by ignoring line segments within bounding boxes. Among other things, this approach can avoid using heavy computing resources (which are required by conventional deep neural networks). In some cases, multiple processors **102** and/or multiple cores or engines of the processor(s) **102** can be used to perform different portions of the process described above.

(41) Although FIG. **2** illustrates one example of a method **200** for vanishing point determination, various changes may be made to FIG. **2**. For example, while shown as a series of steps, various steps in FIG. **2** may overlap, occur in parallel, occur in a different order, or occur any number of times. Although FIG. **3** illustrates one example of a vanishing point determination, various changes may be made to FIG. **3**. For instance, the contents of the image **300** and the results of the vanishing point determination can vary depending on the actual environment around the system **100**.

(42) The vanishing point determined for an image as described above may be used in any suitable manner. For example, in some cases, the vanishing point determined for an image may be used to perform symmetry-based boundary refinement as described below. Note, however, that this use of the vanishing point is for illustration only.

(43) FIG. **4** illustrates an example system **400** supporting symmetry-based boundary refinement according to this disclosure. The system **400** here includes many of the same components described above with respect to the system **100**, and common reference numbers are used in the systems. As shown in FIG. **4**, the processor **102** of FIG. **4** performs the object detection function **108** and the vanishing point detection function **114**. Note that the line segment detection function **110** and the incremental aggregation function **112** are omitted from FIG. **4** for clarity but can be performed by the processor **102** of the system **400** in FIG. **4**.

(44) The processor **102** in FIG. **4** is configured to perform symmetry-based boundary refinement, which generally involves refining one or more bounding boxes or other boundaries for one or more detected objects to more accurately define the boundary of a specified portion of each detected object (such as a front or rear portion of each detected object). Each original bounding box from the object detection function **108** may be defined using the notation $\{\text{right arrow over } (\theta)\}_{\text{sub.bbox}, k} = (x_{\text{sub.k}}, y_{\text{sub.k}}, w_{\text{sub.k}}, h_{\text{sub.k}})$ as described above. Each refined bounding box that is generated as described below may be defined using the notation $\{\text{right arrow over } (\theta)\}_{\text{sub.bbox}, k} = (x_{\text{sub.k}}, y_{\text{sub.k}}, w_{\text{sub.k}}, h_{\text{sub.k}}, \pi_{\text{sub.l}, k}, \pi_{\text{sub.r}, k})$, where $(\pi_{\text{sub.l}, k}, \pi_{\text{sub.r}, k}) \in [0, 1]$ are ratio parameters that represent the refined bounding box (such as when $x_{\text{sub.k}} + w_{\text{sub.k}} \cdot \pi_{\text{sub.l}, k}$ represents the refined box's left boundary edge). Thus, each refined bounding box has a set of $(\pi_{\text{sub.l}, i}, \pi_{\text{sub.r}, i})$ ratios that alter the size of the original bounding box.

(45) In this example, the processor **102** performs a multi-template matching function **402** and a bounding box refinement function **404**. The multi-template matching function **402** generally operates to analyze each bounding box generated by the object detection function **108** using (among other things) the vanishing point identified by the vanishing point detection function **114**. The multi-template matching function **402** uses an expected symmetry of each detected vehicle in an image so that the bounding box refinement function **404** is able to refine the bounding box associated with that detected vehicle. As described in more detail below, the multi-template matching function **402** can analyze different patches on right and left sides of an original bounding box and look for statistically-significant similarities between the patches on the right and left sides of the original bounding box. The similarities between the patches can be used to identify a modification to be applied by the bounding box refinement function

404 to update the original bounding box and generate an updated or refined bounding box. Example operations that can be performed by the multi-template matching function **402** and the bounding box refinement function **404** are described below. One or more refined bounding boxes determined for a scene may be used in any suitable manner, including in the techniques described below.

(46) Note that the functions **402-404** shown in FIG. 4 and described above may be implemented in any suitable manner in the system **400**. For example, in some embodiments, various functions **402-404** may be implemented or supported using one or more software applications or other software instructions that are executed by at least one processor **102**. In other embodiments, at least some of the functions **402-404** can be implemented or supported using dedicated hardware components. In general, the functions **402-404** described above may be performed using any suitable hardware or any suitable combination of hardware and software firmware instructions.

(47) Although FIG. 4 illustrates one example of a system **400** supporting symmetry-based boundary refinement, various changes may be made to FIG. 4. For example, various functions and components shown in FIG. 4 may be combined, further subdivided, replicated, omitted, or rearranged and additional functions and components may be added according to particular needs. Also, as noted above, the functionality for symmetry-based boundary refinement may be used in any other suitable system, and the system may or may not relate to automotive vehicles or other vehicles.

(48) FIG. 5 illustrates an example method **500** for symmetry-based boundary refinement according to this disclosure. The method **500** may, for example, be performed using the components of the system **400** shown in FIG. 4. Note, however, that the method **500** may be performed using any other suitable device or system. Also, during the discussion of the method **500**, reference is made to FIG. 6, which illustrates an example symmetry-based boundary refinement according to this disclosure. The example of the refinement shown in FIG. 6 is for illustration only and is merely meant to illustrate how various steps in FIG. 5 may be performed.

(49) As shown in FIG. 5, a vanishing point is identified using an image at step **502**. This may include, for example, the processor **102** performing the process shown in FIG. 2 to identify a vanishing point {right arrow over (p)}.sub.vanish=(x, y), which can define the x and y coordinates of the determined vanishing point in the image. Left and right virtual ego lines are identified in the image at step **504**. This may include, for example, the processor **102** estimating a left virtual ego line {right arrow over (l)}.sub.ego,l as passing through the determined vanishing point and through the bottom right corner of a bounding box around a nearby vehicle and estimating a right virtual ego line {right arrow over (l)}.sub.ego,r as passing through the determined vanishing point and through the bottom left corner of a bounding box around another nearby vehicle. In some cases, the left virtual ego line can be described as follows:

$$\{\text{right arrow over (l)}\}.\text{sub.ego,l} = a.\text{sub.lx} + b.\text{sub.ly} + c.\text{sub.l} \quad (6)$$

A similar equation may be used for the right virtual ego line. An example of this is shown in FIG. 6, where the image **300** is processed to identify a left virtual ego line **602** and a right virtual ego line **604**. The left and right virtual ego lines and the bottom of the image are used to define a triangular region within the image at step **506**. An example of this is shown in FIG. 6, where the image **300** includes a triangular region **606**. The identified triangular region may be denoted P.sub.triangle.

(50) A bounding box defined for the image is selected at step **508**. This may include, for example, the processor **102** performing the multi-template matching function **402** to select one of the bounding boxes generated by the object detection function **108**. A determination is made whether a defined point of the selected bounding box is inside or outside of the triangular region at step **510**. For convenience, the bottom center point {right arrow over (p)}.sub.cb of the selected bounding box may be used as the defined point, although other points along the selected bounding box may be used as the defined point. This may include, for example, the processor **102** performing the multi-template matching function **402** to determine whether the defined point of the selected bounding box is inside the triangular region, to the left of the triangular region, or to the right of the triangular region. In some embodiments, the processor **102** may perform the following to make this determination:

$$(51) \quad \begin{aligned} &\text{if } \overset{\text{.fwdarw.}}{p}_{\text{cb}}(y) + \frac{a_l}{b_l} \overset{\text{.fwdarw.}}{p}_{\text{cb}}(x) < \frac{c_l}{b_l}, \quad \text{'left'} \\ \overset{\text{.fwdarw.}}{p}_{\text{cb}} \subset P_{\text{triangle}} = \{ &\text{if } \overset{\text{.fwdarw.}}{p}_{\text{cb}}(y) + \frac{a_r}{b_r} \overset{\text{.fwdarw.}}{p}_{\text{cb}}(x) < \frac{c_r}{b_r}, \quad \text{'right'} \\ &\text{else,} \quad \text{'center'} \end{aligned} \quad (7)$$

Here, an object is determined to be within the triangular region (center) or outside the triangular region (left or right). Bounding boxes for objects within the triangular region may be excluded from further processing, while bounding boxes for objects to the left and right of the triangular region may be processed further for refinement.

(52) Assuming the selected bounding box is outside the triangular region, an initial value for the bounding box ratio is defined at step **512**. This may include, for example, the processor **102** performing the multi-template matching function **402** to define the initial bounding box ratio. In some cases, initial bounding box ratio values for

a bounding box may be defined as:

$$(53) \quad \sim_{l,i} = k_l \frac{\overset{\text{.fwdarw.}}{p}_{cb}(x) - \overset{\text{.fwdarw.}}{p}_{cb,l}(x)}{\overset{\text{.fwdarw.}}{p}_{rc}(x) - \overset{\text{.fwdarw.}}{p}_{cb,l}(x)} + (1 - k_l) \frac{\overset{\text{.fwdarw.}}{p}_{cb}(y) - \overset{\text{.fwdarw.}}{p}_{vanish}(y)}{\overset{\text{.fwdarw.}}{p}_{rc}(y) - \overset{\text{.fwdarw.}}{p}_{vanish}(y)}, k_l \in [0, 1] \quad (8)$$

$$\sim_{r,i} = k_r \frac{\overset{\text{.fwdarw.}}{p}_{cb,r}(x) - \overset{\text{.fwdarw.}}{p}_{cb}(x)}{\overset{\text{.fwdarw.}}{p}_{rc}(x) - \overset{\text{.fwdarw.}}{p}_{cb,r}(x)} + (1 - k_r) \frac{\overset{\text{.fwdarw.}}{p}_{cb}(y) - \overset{\text{.fwdarw.}}{p}_{vanish}(y)}{\overset{\text{.fwdarw.}}{p}_{rc}(y) - \overset{\text{.fwdarw.}}{p}_{vanish}(y)}, k_r \in [0, 1] \quad (9) \text{ where:}$$

$$\overset{\text{.fwdarw.}}{p}_{cb,l} = (-\frac{b_l}{a_l} \overset{\text{.fwdarw.}}{p}_{cb}(x) + \frac{c_l}{a_l} \overset{\text{.fwdarw.}}{p}_{cb}(y)) \quad (10) \quad \overset{\text{.fwdarw.}}{p}_{cb,r} = (-\frac{b_r}{a_r} \overset{\text{.fwdarw.}}{p}_{cb}(x) + \frac{c_r}{a_r} \overset{\text{.fwdarw.}}{p}_{cb}(y)) \quad (11)$$

(54) An iterative process occurs to identify and analyze patches from the selected bounding box. Left and right patches in the selected bounding box are selected at step **514**. This may include, for example, the processor **102** performing the multi-template matching function **402** to select a patch of the selected bounding box in a left portion of the selected bounding box and to select a patch of the selected bounding box in a right portion of the selected bounding box. A first of the patches here may be selected randomly, in a predefined manner, or in any other suitable manner, and a second of the patches may be a mirror image of the first patch. In some embodiments, the processor **102** may select a rectangular-shaped patch $\{\text{right arrow over } (\theta)\}.\text{sub.tpl}, k=(x.\text{sub.k}, y.\text{sub.k}, w.\text{sub.k}, h.\text{sub.k})$ from the left edge of the selected bounding box (meaning $x.\text{sub.k}=x.\text{sub.i}+w.\text{sub.i}.\text{Math}.\{\tilde{\text{ over } (\pi)\}.\text{sub.l}, j\}$) and select a mirrored patch $\{\text{right arrow over } (\theta')\}.\text{sub.tpl}, k$ from the residual portion up to the right edge of the selected bounding box (meaning $x.\text{sub.k}+w.\text{sub.k}=x.\text{sub.i}+w.\text{sub.i}.\text{Math}.\{1-\{\tilde{\text{ over } (\pi)\}.\text{sub.r}, i\}\}$).

(55) A measure of similarity between the left and right patches is identified at step **516**. This may include, for example, the processor **102** performing the multi-template matching function **402** to calculate the visual similarity of the image data in the left and right patches. In some embodiments, the similarity can be calculated as a mathematical correlation, such as in the following manner:

$$s.\text{sub.k}, \{\text{right arrow over } (p)\}.\text{sub.match}, k = f(\{\text{right arrow over } (\theta)\}.\text{sub.tpl}, k, \{\text{right arrow over } (\theta')\}.\text{sub.tpl}, k) \quad (12)$$

Here, $s.\text{sub.k}$ denotes the visual similarity between the two patches. The similarity and the patch coordinates (which can be said to represent template coordinates) can be stored at step **518**. This may include, for example, the processor **102** storing a triple of $\{\text{visual similarity } s.\text{sub.k}, \text{left template coordinates } \{\text{right arrow over } (p)\}.\text{sub.left-tpl}=g(\{\text{right arrow over } (p)\}.\text{sub.match}, k, \{\text{right arrow over } (\theta)\}.\text{sub.tpl}, k), \text{right template coordinates } \{\text{right arrow over } (p)\}.\text{sub.right-tpl}=h(\{\text{right arrow over } (p)\}.\text{sub.match}, k, \{\text{right arrow over } (\theta')\}.\text{sub.tpl}, k)\}$. This may be expressed as follows:

$$S.\text{sub.s} \leftarrow \{s.\text{sub.k}, \{\text{right arrow over } (p)\}.\text{sub.left-tpl}, \{\text{right arrow over } (p)\}.\text{sub.right-tpl}\} \quad (13)$$

(56) A determination is made whether to repeat another iteration at step **520**. This may include, for example, the processor **102** determining whether a user-specified number or other number of estimated similarities and template coordinates have been identified. If not, the process can return to step **514**. Otherwise, finalized left and right template coordinates are identified at step **522**. This may include, for example, the processor **102** performing the multi-template matching function **402** to calculate the finalized left and right template coordinates using the stored template coordinates and visual similarities. In some embodiments, the multi-template matching function **402** determines the finalized left and right template coordinates as a weighted combination of the stored left and right template coordinates, where the visual similarities are used as the weights for the combination. This can be expressed as:

$$(57) \quad \overset{\text{.fwdarw.}}{p}_{\text{left-tpl}} = \text{.Math.}_{k=1, \text{.Math.}, K} \tilde{S}_k \cdot \overset{\text{.fwdarw.}}{p}_{\text{left-tpl}}, \tilde{S}_k = \frac{S_k}{\text{.Math.}_{j=1, \text{.Math.}, K} S_j} \quad (14)$$

$$\overset{\text{.fwdarw.}}{p}_{\text{right-tpl}} = \text{.Math.}_{k=1, \text{.Math.}, K} \tilde{S}_k \cdot \overset{\text{.fwdarw.}}{p}_{\text{right-tpl}}, \tilde{S}_k = \frac{S_k}{\text{.Math.}_{j=1, \text{.Math.}, K} S_j} \quad (15)$$

(58) The boundary ratio for the selected bounding box is updated and used to refine the selected bounding box at step **524**. This may include, for example, the processor **102** performing the multi-template matching function **402** to update the boundary ratio. In some embodiments, this can be expressed as follows:

$$(59) \quad -_{l,i} = \frac{\overset{\text{.fwdarw.}}{p}_{\text{left-tpl}} - x_i}{w_i} \quad (16) \quad -_{r,i} = \frac{x_i + w_i - \overset{\text{.fwdarw.}}{p}_{\text{right-tpl}}}{w_i} \quad (17)$$

This may also include the processor **102** performing the bounding box refinement function **404** to modify the current selected bounding box based on the updated boundary ratio. For instance, the processor **102** can refine the selected bounding box so that it has the updated boundary ratios within the image.

(60) A determination is made whether one or more additional bounding boxes need to be processed at step **526**. If so, the process can return to step **508** to select and refine another bounding box. Otherwise, the process can end,

and the refined bounding box(es) can be used in any suitable manner.

(61) Example results obtained using the process are shown in FIG. 6, where the bounding box **302** is refined to a bounding box **302'** and the bounding box **304** is refined to a bounding box **304'**. The bounding box **302'** here more accurately defines the boundary of the rear portion of one vehicle, and the bounding box **304'** here more accurately defines the boundary of the rear portion of another vehicle. Essentially, the bounding box **302'** identifies a generally symmetrical portion of the left vehicle, and the bounding box **304'** identifies a generally symmetrical portion of the right vehicle. The ability to accurately identify the rear portions of nearby vehicles may be used in various ways, such as to identify a direction of travel or a change in the direction of travel of the left or right vehicle or to identify a surface of another vehicle to be used for depth estimation.

(62) Although FIG. 5 illustrates one example of a method **500** for symmetry-based boundary refinement, various changes may be made to FIG. 5. For example, while shown as a series of steps, various steps in FIG. 5 may overlap, occur in parallel, occur in a different order, or occur any number of times. Although FIG. 6 illustrates one example of a symmetry-based boundary refinement, various changes may be made to FIG. 6. For instance, the contents of the image **600** and the results of the symmetry-based boundary refinement can vary depending on the actual environment around the system **100**.

(63) The refined bounding boxes or other boundaries determined for an image as described above may be used in any suitable manner. For example, in some cases, the refined bounding boxes determined for an image may be used to perform component detection as described below. Note, however, that this use of the refined bounding boxes is for illustration only.

(64) FIG. 7 illustrates an example system **700** supporting component detection according to this disclosure. The system **700** here includes many of the same components described above with respect to the systems **100** and **400**, and common reference numbers are used in the systems. As shown in FIG. 7, the processor **102** of FIG. 7 performs the object detection function **108** and the bounding box refinement function **404**. Note that the line segment detection function **110**, incremental aggregation function **112**, vanishing point detection function **114**, and multi-template matching function **402** are omitted from FIG. 7 for clarity but can be performed by the processor **102** of the system **700** in FIG. 7.

(65) The processor **102** in FIG. 7 is configured to perform component detection in which one or more individual components of other vehicles (such as license plates or taillights) are identified. Here, the processor **102** performs an integral image generation function **702**, which generally operates to produce an integral image using the image from a camera **104a**. In some embodiments, an integral image $I(\cdot, \cdot)$ can be generated as follows:

$$I(x,y)=\sum_{x'\leq x,y'\leq y}i(x',y') \quad (18)$$

$$I(x,y)=i(x,y)+I(x-1,y)+I(x,y-1)-I(x-1,y-1) \quad (19)$$

where:

$$I(x',y')=0, x',y'<0 \quad (20)$$

(66) The integral image from the integral image generation function **702** and one or more refined bounding boxes from the bounding box refinement function **404** are provided to a fast block comparison function **704**. The fast block comparison function **704** generally operates to select different (mirrored) regions within each refined bounding box and analyze the regions to compute probabilistic similarities between the regions, which can be based on the integral image. As described in more detail below, the fast block comparison function **704** can analyze different mirrored regions of a refined bounding box and look for regions that are probabilistically similar. The similarities between the regions can be used by a component detection function **706** to identify regions in the refined bounding box associated with one or more specific components of a vehicle. Example operations that can be performed by the fast block comparison function **704** and the component detection function **706** are described below. One or more identified vehicle components may be used in any suitable manner.

(67) Note that the functions **702-706** shown in FIG. 7 and described above may be implemented in any suitable manner in the system **700**. For example, in some embodiments, various functions **702-706** may be implemented or supported using one or more software applications or other software instructions that are executed by at least one processor **102**. In other embodiments, at least some of the functions **702-706** can be implemented or supported using dedicated hardware components. In general, the functions **702-706** described above may be performed using any suitable hardware or any suitable combination of hardware and software/firmware instructions.

(68) Although FIG. 7 illustrates one example of a system **700** supporting component detection, various changes may be made to FIG. 7. For example, various functions and components shown in FIG. 7 may be combined, further subdivided, replicated, omitted, or rearranged and additional functions and components may be added according to particular needs. Also, as noted above, the functionality for component detection may be used in any other suitable system, and the system may or may not relate to automotive vehicles or other vehicles.

(69) FIG. 8 illustrates an example method **800** for component detection according to this disclosure. The method **800** may, for example, be performed using the components of the system **700** shown in FIG. 7. Note, however, that

the method **800** may be performed using any other suitable device or system. Also, during the discussion of the method **800**, reference is made to FIG. **9**, which illustrates an example component detection according to this disclosure. The example of the detection shown in FIG. **9** is for illustration only and is merely meant to illustrate how various steps in FIG. **8** may be performed.

(70) As shown in FIG. **8**, one or more refined bounding boxes for one or more objects in an image are identified at step **802**. This may include, for example, the processor **102** performing the process shown in FIG. **2** to identify a vanishing point and performing the process shown in FIG. **5** to generate the refined bounding box(es). An integral image of a scene is generated at step **804**. This may include, for example, processor **102** generating the integral image as described above.

(71) A refined bounding box is selected at step **806**. This may include, for example, the processor **102** performing the fast block comparison function **704** to select one of the refined bounding boxes generated by the bounding box refinement function **404**. Mirrored regions in the selected refined bounding box are selected at step **808**. This may include, for example, the processor **102** performing the fast block comparison function **704** to select mirrored regions in the selected refined bounding box randomly, in a predefined manner, or in any other suitable manner. In some embodiments, the processor **102** may select a rectangular-shaped region $\{\text{right arrow over } (\theta)\}.\text{sub.left},k = (x.\text{sub.k}, y.\text{sub.k}, w.\text{sub.k}, h.\text{sub.k})$ inside the selected refined boundary box and another rectangular region $\{\text{right arrow over } (\theta)\}.\text{sub.right},k$ inside the selected refined boundary box (but mirrored against a middle line of the refined boundary box).

(72) A probabilistic distribution of each region is modeled at step **810**. This may include, for example, the processor **102** performing the fast block comparison function **704** to model the probabilistic distribution of each rectangular region or other defined region. For simplicity, a normal distribution parameterization $\text{custom character.sub.k}(\mu.\text{sub.k}, \sigma.\text{sub.k})$ may be selected, where a mean $\mu.\text{sub.k}$ and a standard deviation $\sigma.\text{sub.k}$ are computed using the integral image I and $I'(\cdot, \cdot) = \Sigma.\text{sub.x}' \leq x, y' \leq y i(x', y') .\text{sup.2}$. This can be expressed as:

$$(73) \quad k = \frac{I(\text{right arrow over } (\theta).\text{sub.},k)}{N_k}, \quad k = \frac{\sqrt{I'(\text{right arrow over } (\theta).\text{sub.},k)}}{N_k} \quad (21)$$

Here, $N.\text{sub.k}$ represents a number of pixels in a region $I(\{\text{right arrow over } (\theta)\}.\text{sub.},k)$.

(74) A probabilistic similarity between the two regions is determined at step **812**. This may include, for example, the processor **102** performing the fast block comparison function **704** to calculate the probabilistic similarity between the distributions of the two regions calculated in step **810**. In some embodiments, the probabilistic similarity between two rectangular regions can be determined via their normal distribution forms

custom character custom character . In particular embodiments, a variant of the Kullback-Leibler divergence (KLD), referred to as the Jensen Shannon divergence (JSD), may be used to overcome the asymmetry associated with KLD. In these embodiments, the probabilistic similarity may be determined as follows:

$$(75) \quad D_{JS,k} = \frac{D_{KL,k}(\text{left}.\text{Math}.\text{right}) + D_{KL,k}(\text{right}.\text{Math}.\text{left})}{2} \quad (22) \quad \text{where:}$$

$$D_{KL,k}(\text{left}.\text{Math}.\text{right}) = \log \frac{\text{right}}{\text{left}} + \frac{\frac{2}{\text{left}} + (\frac{\text{left}}{2} - \frac{2}{\text{right}})}{2} - \frac{1}{2} \quad (23)$$

Note, however, that other measures of probabilistic similarity may be used here. The probabilistic similarity and the coordinates of the two regions can be stored at step **814**. This may include, for example, the processor **102** storing a tuple of {JSD value $D.\text{sub.JS},k$, left rectangular region $\{\text{right arrow over } (\theta)\}.\text{sub.left},k$, right rectangular region $\{\text{right arrow over } (\theta)\}.\text{sub.right},k$ }. This may be expressed as follows:

$$S.\text{sub.p} \leftarrow \{D.\text{sub.JS},k, \{\text{right arrow over } (\theta)\}.\text{sub.left},k, \{\text{right arrow over } (\theta)\}.\text{sub.right},k\} \quad (24)$$

(76) A determination is made whether to repeat another iteration at step **816**. This may include, for example, the processor **102** determining whether a user-specified number or other number of estimated similarities and regions have been identified. If not, the process can return to step **808**. Otherwise, finalized mirrored regions associated with the selected refined bounding box are identified at step **818**. This may include, for example, the processor **102** performing the component detection function **706** to determine the finalized mirrored regions using the stored region coordinates and probabilistic similarities. In some embodiments, the component detection function **706** determines the finalized mirrored regions as a weighted combination of the stored region coordinates, where the probabilistic similarities are used as the weights for the combination. This can be expressed as:

$$(77) \quad 0.\text{sub.left},k = \text{Math}.\text{sub.k} = 1, \text{Math}.\text{sub.k} \quad \text{left},k \quad \text{Math}.\tilde{D}_{JS,k}, \tilde{D}_{JS,k} = \frac{D_{JS,k}}{\text{Math}.\text{sub.j} = 1, \text{Math}.\text{sub.K} D_{JS,j}} \quad (25)$$

$$\text{right},k = \text{Math}.\text{sub.k} = 1, \text{Math}.\text{sub.k} \quad \text{right},k \quad \text{Math}.\tilde{D}_{JS,k}, \tilde{D}_{JS,k} = \frac{D_{JS,k}}{\text{Math}.\text{sub.j} = 1, \text{Math}.\text{sub.K} D_{JS,j}} \quad (26)$$

(78) A determination is made whether one or more additional bounding boxes need to be processed at step **820**. If so, the process can return to step **806** to select and refine another bounding box. Otherwise, the process can end,

and the identified vehicle component(s) can be used in any suitable manner.

(79) An example of this is shown in FIG. 9, where the image 300 is processed to identify two regions 902 and 904 within the refined bounding box 304'. As can be seen here, the processor 102 has identified the two regions 902 and 904 of the image 300 associated with the taillights of another vehicle since, based on the process described above, these two regions 902 and 904 in the refined bounding box 304' are probabilistically similar. The same type of process may be used, for example, to identify two regions containing the license plate of the other vehicle. The positions of these components of the other vehicle may be used for various purposes, such as identifying a direction of travel or a change in the direction of travel of the other vehicle relative to the target vehicle or identifying a surface of the other vehicle to be used for depth estimation.

(80) Although FIG. 8 illustrates one example of a method 800 for component detection, various changes may be made to FIG. 8. For example, while shown as a series of steps, various steps in FIG. 8 may overlap, occur in parallel, occur in a different order, or occur any number of times. Although FIG. 9 illustrates one example of a component detection, various changes may be made to FIG. 9. For instance, the contents of the image 900 and the results of the component detection can vary depending on the actual environment around the system 100.

(81) FIG. 10 illustrates an example usage of vanishing point determination, symmetry-based boundary refinement, and/or component detection according to this disclosure. A system 1000 here includes many of the same components described above with respect to the systems 1000, 400, and 700, and common reference numbers are used in the systems. As shown in FIG. 10, the processor 102 of FIG. 10 performs the vanishing point detection function 114, the bounding box refinement function 404, and the component detection function 706. Note that the other functions described above are omitted from FIG. 10 for clarity but can be performed by the processor 102 of the system 1000 in FIG. 10.

(82) The processor 102 performs a decision planning function 1002, which generally uses one or more determined vanishing points, one or more refined bounding boxes, and/or one or more identified vehicle components to determine how to adjust the operation of the system 1000. For example, in an automotive vehicle, the decision planning function 1002 may determine whether (and how) to change the steering direction of the vehicle, whether (and how) to apply the brakes or accelerate the vehicle, or whether (and how) to trigger an audible, visible, haptic, or other warning. The warning may indicate that the system 1000 is near another vehicle, obstacle, or person, is departing from a current lane in which the vehicle is traveling, or is approaching a possible impact location with another vehicle, Obstacle, or person. In general, the identified adjustments determined by the decision planning function 1002 can vary widely based on the specific application.

(83) The decision planning function 1002 can interact with one or more control functions 1004, each of which can be used to adjust or control the operation of one or more actuators 1006 in the system 1000. For example, in an automotive vehicle, the one or more actuators 1006 may represent one or more brakes, electric motors, or steering components of the vehicle, and the control function(s) 1004 can be used to apply or discontinue application of the brakes, speed up or slow down the electric motors, or change the steering direction of the vehicle. In general, the specific way(s) in which detected objects can be used may vary depending on the specific system 1000 in which object detection is being used.

(84) Note that the functions 1002-1006 shown in FIG. 10 and described above may be implemented in any suitable manner in the system 100. For example, in some embodiments, various functions 1002-1006 may be implemented or supported using one or more software applications or other software instructions that are executed by at least one processor 102. In other embodiments, at least some of the functions 1002-1006 can be implemented or supported using dedicated hardware components. In general, the functions 1002-1006 described above may be performed using any suitable hardware or any suitable combination of hardware and software/firmware instructions.

(85) Although FIG. 10 illustrates one example usage of vanishing point determination, symmetry-based boundary refinement, and/or component detection, various changes may be made to FIG. 10. For example, various functions and components shown in FIG. 10 may be combined, further subdivided, replicated, omitted, or rearranged and additional functions and components may be added according to particular needs. Also, while vanishing point determination, symmetry-based boundary refinement, and component detection are used here, the system 1000 may use a single one of these functions or any desired combination of these functions.

(86) Note that many functional aspects of the embodiment described above can be implemented using any suitable hardware or any suitable combination of hardware and software/firmware instructions. In some embodiments, at least some functional aspects of the embodiments described above can be embodied as software instructions that are executed by one or more unitary or multi-core central processing units or other processing device(s). In other embodiments at least some functional aspects of the embodiments described above can be embodied using one or more application specific integrated circuits (ASICs). When implemented using one or more ASICs, any suitable integrated circuit design and manufacturing techniques may be used, such as those that can be automated using electronic design automation (EDA) tools. Examples of such tools include tools provided by SYNOPSYS, INC.,

(87) FIG. **11** illustrates an example design flow **1100** for employing one or more tools to design hardware that implements one or more control functions according to this disclosure. More specifically, the design flow **1100** here represents a simplified ASIC design flow employing one or more EDA tools for designing and facilitating fabrication of ASICs that implement at least some functional aspects of the various embodiments described above.

(88) As shown in FIG. **11**, a functional design of an ASIC is created at step **1102**. For any portion of the ASIC design that is digital in nature, in some cases, this may include expressing the digital functional design by generating register transfer level (RTL) code in a hardware descriptive language (HDL), such as VEIDL or VERILOG. A functional verification (such as a behavioral simulation) can be performed on HDL data structures to ensure that the RTL code that has been generated is in accordance with logic specifications. In other cases, a schematic of digital logic can be captured and used, such as through the use of a schematic capture program. For any portion of the ASIC design that is analog in nature, this may include expressing the analog functional design by generating a schematic, such as through the use of a schematic capture program. The output of the schematic capture program can be converted (synthesized), such as into gate/transistor level netlist data structures. Data structures or other aspects of the functional design are simulated, such as by using a simulation program with integrated circuits emphasis (SPICE), at step **1104**. This may include, for example, using the SPICE simulations or other simulations to verify that the functional design of the ASIC performs as expected.

(89) A physical design of the ASIC is created based on the validated data structures and other aspects of the functional design at step **1106**. This may include, for example, instantiating the validated data structures with their geometric representations. In some embodiments, creating a physical layout includes “floor-planning,” where gross regions of an integrated circuit chip are assigned and input/output (I/O) pins are defined. Also, hard cores (such as arrays, analog blocks, inductors, etc.) can be placed within the gross regions based on design constraints (such as trace lengths, timing, etc.). Clock wiring, which is commonly referred to or implemented as clock trees, can be placed within the integrated circuit chip, and connections between gates/analog blocks can be routed within the integrated circuit chip. When all elements have been placed, a global and detailed routing can be performed to connect all of the elements together. Post-wiring optimization may be performed to improve performance (such as timing closure), noise (such as signal integrity), and yield. The physical layout can also be modified where possible while maintaining compliance with design rules that are set by a captive, external, or other semiconductor manufacturing foundry of choice, which can make the ASIC more efficient to produce in bulk. Example modifications may include adding extra vias or dummy metal/diffusion/poly layers.

(90) The physical design is verified at step **1108**. This may include, for example, performing design rule checking (DRC) to determine whether the physical layout of the ASIC satisfies a series of recommended parameters, such as design rules of the foundry. In some cases, the design rules represent a series of parameters provided by the foundry that are specific to a particular semiconductor manufacturing process. As particular examples, the design rules may specify certain geometric and connectivity restrictions to ensure sufficient margins to account for variability in semiconductor manufacturing processes or to ensure that the ASICs work correctly. Also, in some cases, a layout versus schematic (LVS) check can be performed to verify that the physical layout corresponds to the original schematic or circuit diagram of the design. In addition, a complete simulation may be performed to ensure that the physical layout phase is properly done.

(91) After the physical layout is verified, mask generation design data is generated at step **1110**. This may include, for example, generating mask generation design data for use in creating photomasks to be used during ASIC fabrication. The mask generation design data may have any suitable form, such as GDSII data structures. This step may be said to represent a “tape-out” for preparation of the photomasks. The GDSII data structures or other mask generation design data can be transferred through a communications medium (such as via a storage device or over a network) from a circuit designer or other party to a photomask supplier/maker or to the semiconductor foundry itself. The photomasks can be created and used to fabricate ASIC devices at step **1112**.

(92) Although FIG. **11** illustrates one example of a design flow **1100** for employing one or more tools to design hardware that implements one or more vehicle control functions, various changes may be made to FIG. **11**. For example, at least some functional aspects of the various embodiments described above may be implemented in any other suitable manner.

(93) FIG. **12** illustrates an example device **1200** supporting execution of one or more tools to design hardware that implements one or more vehicle control functions according to this disclosure. The device **1200** may, for example, be used to implement at least part of the design flow **1100** shown in FIG. **11**. However, the design flow **1100** may be implemented in any other suitable manner.

(94) As shown in FIG. **12**, the device **1200** denotes a computing device or system that includes at least one processing device **1202**, at least one storage device **1204**, at least one communications unit **1206**, and at least one input/output (I/O) unit **1208**. The processing device **1202** may execute instructions that can be loaded into a memory **1210**. The processing device **1202** includes any suitable number(s) and type(s) of processors or other

processing devices in any suitable arrangement. Example types of processing devices **1202** include one or more microprocessors, microcontrollers, digital signal processors (DSPs), application specific integrated circuits (ASICs), field programmable gate arrays (FPGAs), or discrete circuitry.

(95) The memory **1210** and a persistent storage **1212** are examples of storage devices **1204**, which represent any structure(s) capable of storing and facilitating retrieval of information (such as data, program code, and/or other suitable information on a temporary or permanent basis). The memory **1210** may represent a random access memory or any other suitable volatile or non-volatile storage device(s). The persistent storage **1212** may contain one or more components or devices supporting longer-term storage of data, such as a read only memory, hard drive, Flash memory, or optical disc.

(96) The communications unit **1206** supports communications with other systems or devices. For example, the communications unit **1206** can include a network interface card or a wireless transceiver facilitating communications over a wired or wireless network. The communications unit **1206** may support communications through any suitable physical or wireless communication link(s).

(97) The I/O unit **1208** allows for input and output of data. For example, the I/O unit **1208** may provide a connection for user input through a keyboard, mouse, keypad, touchscreen, or other suitable input device. The I/O unit **1208** may also send output to a display or other suitable output device. Note, however, that the I/O unit **1208** may be omitted if the device **1200** does not require local I/O, such as when the device **1200** represents a server or other device that can be accessed remotely.

(98) The instructions that are executed by the processing device **1202** include instructions that implement at least part of the design flow **1100**. For example, the instructions that are executed by the processing device **1202** may cause the processing device **1202** to generate or otherwise obtain functional designs, perform simulations, generate physical designs, verify physical designs, perform tape-outs, or create/use photomasks (or any combination of these functions). As a result, the instructions that are executed by the processing device **1202** support the design and fabrication of ASIC devices or other devices that implement one or more vehicle control functions described above.

(99) Although FIG. **12** illustrates one example of a device **1200** supporting execution of one or more tools to design hardware that implements one or more vehicle control functions, various changes may be made to FIG. **12**. For example, computing and communication devices and systems come in a wide variety of configurations, and FIG. **12** does not limit this disclosure to any particular computing or communication device or system.

(100) Note that while the various functions described above are often described as using or being based on images captured using one or more cameras **104a**, other types of sensors **104** may be used to provide sensor data for use by these functions. For example, abstracted image representations or pseudo-image representations of RADAR or LIDAR measurements may be processed as described above to identify vanishing points, perform symmetry-based boundary refinement, and/or engage in component detection.

(101) Also note that while bounding boxes are often described above as being generated, refined, processed, or otherwise used to represent the boundaries of objects, these boundaries may be expressed in any other suitable manner. Example types of boundaries that may be defined and used in the various systems, techniques, and processes described above include two-dimensional bounding boxes, three-dimensional cuboids, and/or boundaries around recognized objects (such as bridges, buildings, trees, light poles or other poles, traffic signs, traffic lights, etc.).

(102) In some embodiments, various functions described in this patent document are implemented or supported using machine-readable instructions that are stored on a non-transitory machine-readable medium. The phrase “machine-readable instructions” includes any type of instructions, including source code, object code, and executable code. The phrase “non-transitory machine-readable medium” includes any type of medium capable of being accessed by one or more processing devices or other devices, such as a read only memory (ROM), a random access memory (RAM), a Flash memory, a hard disk drive (HDD), or any other type of memory. A “non-transitory” medium excludes wired, wireless, optical, or other communication links that transport transitory electrical or other signals. Non-transitory media include media where data can be permanently stored and media where data can be stored and later overwritten.

(103) It may be advantageous to set forth definitions of certain words and phrases used throughout this patent document. The terms “include” and “comprise,” as well as derivatives thereof, mean inclusion without limitation. The term “or” is inclusive, meaning and/or. The phrase “associated with,” as well as derivatives thereof, may mean to include, be included within, interconnect with, contain, be contained within, connect to or with, couple to or with, be communicable with, cooperate with, interleave, juxtapose, be proximate to, be bound to or with, have, have a property of, have a relationship to or with, or the like. The phrase “at least one of,” when used with a list of items, means that different combinations of one or more of the listed items may be used, and only one item in the list may be needed. For example, “at least one of: A, B, and C” includes any of the following combinations: A, B, C, A and B, A and C, B and C, and A and B and C.

(104) The description in the present application should not be read as implying that any particular element, step, or function is an essential or critical element that must be included in the claim scope. The scope of patented subject matter is defined only by the allowed claims. Moreover, none of the claims invokes 35 U.S.C. § 112(f) with respect to any of the appended claims or claim elements unless the exact words “means for” or “step for” are explicitly used in the particular claim, followed by a participle phrase identifying a function. Use of terms such as (but not limited to) “mechanism,” “module,” “device,” “unit,” “component,” “element,” “member,” “apparatus,” “machine,” “system,” “processor,” or “controller” within a claim is understood and intended to refer to structures known to those skilled in the relevant art, as further modified or enhanced by the features of the claims themselves, and is not intended to invoke 35 U.S.C. § 112(f).

(105) While this disclosure has described certain embodiments and generally associated methods, alterations and permutations of these embodiments and methods will be apparent to those skilled in the art. Accordingly, the above description of example embodiments does not define or constrain this disclosure. Other changes, substitutions, and alterations are also possible without departing from the spirit and scope of this disclosure, as defined by the following claims.

Claims

1. A method comprising: obtaining, using at least one processing device, a vanishing point and a boundary based on image data associated with a scene, the boundary associated with a detected object within the scene; repeatedly, during multiple iterations and using the at least one processing device, (i) identifying multiple patches within the boundary and (ii) determining a similarity of the image data contained within the multiple patches; identifying, using the at least one processing device, a modification to be applied to the boundary based on the identified patches and the determined similarities; and generating, using the at least one processing device, a refined boundary based on the modification, the refined boundary identifying a specified portion of the detected object, wherein, during each iteration, the multiple patches within the boundary include a first patch and a mirrored second patch within the boundary.
2. The method of claim 1, further comprising: identifying virtual left and right ego lines using the image data, each ego line passing through the vanishing point; defining a triangular region based on the virtual left and right ego lines; and selecting the boundary as having a specified point outside the triangular region.
3. The method of claim 1, further comprising: identifying an initial boundary ratio; and updating the initial boundary ratio to generate an updated boundary ratio based on the modification, wherein the refined boundary is generated based on the updated boundary ratio.
4. The method of claim 1, wherein: each iteration generates estimated template coordinates based on the identified patches within the boundary; and identifying the modification comprises generating a weighted combination of the estimated template coordinates while using the determined similarities as weights for the estimated template coordinates, the modification representing the weighted combination.
5. The method of claim 1, further comprising: determining a position of the vanishing point based on multiple collections of line segments identified based on the image data.
6. The method of claim 1, wherein: the detected object comprises a vehicle; and the specified portion of the detected object comprises a rear portion of the vehicle.
7. The method of claim 1, further comprising: identifying one or more components of the detected object based on the refined boundary.
8. The method of claim 1, further comprising: identifying at least one action to be performed based on the refined boundary; and performing the at least one action.
9. The method of claim 8, wherein the at least one action comprises at least one of: an adjustment to at least one of: a steering of a vehicle, a speed of the vehicle, an acceleration of the vehicle, and a braking of the vehicle; and an activation of an audible, visible, or haptic warning.
10. An apparatus comprising: at least one processing device configured to: obtain a vanishing point and a boundary based on image data associated with a scene, the boundary associated with a detected object within the scene; repeatedly, during multiple iterations, (i) identify multiple patches within the boundary and (ii) determine a similarity of the image data contained within the multiple patches; identify a modification to be applied to the boundary based on the identified patches and the determined similarities; and generate a refined boundary based on the modification, the refined boundary identifying a specified portion of the detected object, wherein, during each iteration, the multiple patches within the boundary include a first patch and a mirrored second patch within the boundary.
11. The apparatus of claim 10, wherein the at least one processing device is further configured to: identify virtual left and right ego lines using the image data, each ego line passing through the vanishing point; define a triangular

region based on the virtual left and right ego lines; and select the boundary as having a specified point outside the triangular region.

12. The apparatus of claim 10, wherein: the at least one processing device is further configured to: identify an initial boundary ratio; and update the initial boundary ratio to generate an updated boundary ratio based on the modification; and the at least one processing device is configured to generate the refined boundary based on the updated boundary ratio.

13. The apparatus of claim 10, wherein: the at least one processing device is configured, during each iteration, to generate estimated template coordinates based on the identified patches within the boundary; and to identify the modification, the at least one processing device is configured to generate a weighted combination of the estimated template coordinates while using the determined similarities as weights for the estimated template coordinates, the modification representing the weighted combination.

14. The apparatus of claim 10, wherein the at least one processing device is further configured to determine a position of the vanishing point based on multiple collections of line segments identified based on the image data.

15. The apparatus of claim 10, wherein: the detected object comprises a vehicle; and the specified portion of the detected object comprises a rear portion of the vehicle.

16. The apparatus of claim 10, wherein the at least one processing device is further configured to identify one or more components of the detected object based on the refined boundary.

17. The apparatus of claim 10, wherein the at least one processing device is further configured to: identify at least one action to be performed based on the refined boundary; and perform the at least one action.

18. The apparatus of claim 17, wherein the at least one action comprises at least one of: an adjustment to at least one of: a steering of a vehicle, a speed of the vehicle, an acceleration of the vehicle, and a braking of the vehicle; and an activation of an audible, visible, or haptic warning.

19. A non-transitory machine-readable medium containing instructions that when executed cause at least one processor to: obtain a vanishing point and a boundary based on image data associated with a scene, the boundary associated with a detected object within the scene; repeatedly, during multiple iterations, (i) identify multiple patches within the boundary and (ii) determine a similarity of the image data contained within the multiple patches; identify a modification to be applied to the boundary based on the identified patches and the determined similarities; and generate a refined boundary based on the modification, the refined boundary identifying a specified portion of the detected object, wherein, during each iteration, the multiple patches within the boundary include a first patch and a mirrored second patch within the boundary.

20. The non-transitory machine-readable medium of claim 19, further containing instructions that when executed cause the at least one processor to: identify virtual left and right ego lines using the image data, each ego line passing through the vanishing point; define a triangular region based on the virtual left and right ego lines; and select the boundary as having a specified point outside the triangular region.

21. The non-transitory machine-readable medium of claim 19, wherein: the medium further contains instructions that when executed cause the at least one processor to: identify an initial boundary ratio; and update the initial boundary ratio to generate an updated boundary ratio based on the modification; and the instructions that when executed cause the at least one processor to generate the refined boundary comprise instructions that when executed cause the at least one processor to generate the refined boundary based on the updated boundary ratio.

22. The non-transitory machine-readable medium of claim 19, wherein: the instructions when executed cause the at least one processor, during each iteration, to generate estimated template coordinates based on the identified patches within the boundary; and the instructions that when executed cause the at least one processor to identify the modification comprise instructions that when executed cause the at least one processor to generate a weighted combination of the estimated template coordinates while using the determined similarities as weights for the estimated template coordinates, the modification representing the weighted combination.

23. The non-transitory machine-readable medium of claim 19, further containing instructions that when executed cause the at least one processor to determine a position of the vanishing point based on multiple collections of line segments identifying based on the image data.

24. The non-transitory machine-readable medium of claim 19, wherein: the detected object comprises a vehicle; and the specified portion of the detected object comprises a rear portion of the vehicle.

25. The non-transitory machine-readable medium of claim 19, further containing instructions that when executed cause the at least one processor to identify one or more components of the detected object based on the refined boundary.

26. The non-transitory machine-readable medium of claim 19, further containing instructions that when executed cause the at least one processor to: identify at least one action to be performed based on the refined boundary; and perform the at least one action.

27. The non-transitory machine-readable medium of claim 26, wherein the at least one action comprises at least

one of: an adjustment to at least one of: a steering of a vehicle, a speed of the vehicle, an acceleration of the vehicle, and a braking of the vehicle; and an activation of an audible, visible, or haptic warning.
