

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12418482
Kind Code	B1
Date of Patent	September 16, 2025
Inventor(s)	Vij; Shivansh et al.

Egress migration of virtualized software services

Abstract

Provided herein are various enhancements for low downtime migration of virtualized software services across different instantiations, which may include real-time migration over different cloud/server providers and platforms, physical locations, network locations, and across different network elements. Examples herein include handling of migration of data and state for virtualized software services, migration of ingress and egress traffic for the software services, and migration of other various operations aspects applicable to virtualized software services. In many instances, a client node retains the same network addressing used to reach the virtualized software services even as the virtualized software services move to different network locations and physical locations.

Inventors:	Vij; Shivansh (Vancouver, CA), Sørlie; Alex (Bergen, NO)
Applicant:	Loophole Labs, Inc. (New York, NY)
Family ID:	1000008477787
Assignee:	Loophole Labs, Inc. (New York, NY)
Appl. No.:	19/081676
Filed:	March 17, 2025

Related U.S. Application Data

continuation parent-doc US 18926569 20241025 US 12255818 child-doc US 19081676
us-provisional-application US 63671933 20240716

Publication Classification

Int. Cl.: H04L45/741 (20220101); H04L45/586 (20220101); H04L61/251 (20220101)

U.S. Cl.:

CPC **H04L45/741** (20130101); **H04L45/586** (20130101); **H04L61/251** (20130101);

Field of Classification Search

CPC: H04L (45/741); H04L (45/586); H04L (61/251)

USPC: 709/238

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2018/0213036	12/2017	Vasetsky	N/A	H04L 41/082
2023/0062733	12/2022	Chu	N/A	G06F 11/0712
2023/0185604	12/2022	Chakaravarthy	718/104	H04L 67/10
2024/0028365	12/2023	Li	N/A	G06F 9/542

Primary Examiner: Asres; Hermon

Background/Summary

RELATED APPLICATIONS (1) This application is a continuation of, and claims priority to, U.S. patent application Ser. No. 18/926,569, entitled “REDUCED DOWNTIME MIGRATION OF VIRTUALIZED SOFTWARE SERVICES,” and filed Oct. 25, 2024. This application claims the benefit of and priority to U.S. Provisional Patent Application 63/671,933, titled “REDUCED DOWNTIME MIGRATION OF VIRTUALIZED SOFTWARE SERVICES,” filed Jul. 16, 2024, which is hereby incorporated by reference in its entirety.

TECHNICAL BACKGROUND

(1) Virtualization of computing environments provides for the abstraction of physical hardware resources and subsequent allocation of such resources among concurrent virtual machines or other compartmentalized software structures. A hypervisor is typically employed as a management layer between physical hardware and the virtual machines, and can perform resource scheduling, resource sharing, resource allocation, hardware abstraction, among other functions. Among these functions are providing primary/main memory (e.g., random access memory or RAM) and secondary storage (e.g., disk or solid-state storage) emulation to the virtual machines. A virtual machine can thus be provided with an appropriate bundle of emulated or virtualized hardware resources, which can be deployed by the hypervisor on-the-fly as these virtual machines are instantiated and de-instantiated on an associated underlying set of physical hardware.

(2) Cloud and distributed computing schemes have enabled deployment of computing resources across various physical data centers, cloud computing centers, and other physical locations, which can see dynamic deployment of virtual machines which change in quantity and location in accordance with available computing/network resources, costs, demand, bandwidth availability, desired qualities of service, and other various factors. A common use case is having a client node, such as an end user, requesting access to a remotely located computing resource over a network link, with a software service responsively spawned as a virtual machine to service that client node.

(3) As operating conditions change during an active session supporting a software service, a migration of a virtual machine to another set of hardware might be desired. Operating conditions include costs of using corresponding computing resources, network conditions (e.g., availability, bandwidth, latency), workload patterns, resource utilization capabilities, time of day, and other conditions. The migration can include “live” migration techniques that attempt to migrate a currently-active software service without interruption of service to a client node.

(4) However, current migration techniques still produce lag and interruptions to an end user or client node, and can require relatively long periods to copy data over a network link to instantiate new virtual machines, such as when moved to a different data center. Moreover, migrating network addressing and network routing resources associated with computing resources can be complex and force client nodes to halt existing network connections and establish new network links for migrated resources. This can increase client node complexity, as well as introduce unwanted opportunities for downtime, service latency, service interruptions, dropped/lost network packets, and overall reduced quality of end user experiences.

SUMMARY

(5) Provided herein are various enhancements for low downtime migration of virtualized software services across different instantiations, which may include real-time migration over different cloud/server providers and platforms, physical locations, network locations, and across different network elements. Migration of virtualized software services can involve moving a virtual machine from one host to another, with the intention of reducing disrupted service by minimizing downtime during the process. Examples herein include enhanced handling of migration of data and state for virtualized software services, migration of network arrangements for ingress and egress traffic for the software services, and migration of other various operations aspects applicable to virtualized software services. In many instances, a client node retains the same network addressing used to reach the virtualized software services even as the virtualized software services move to different network locations and physical locations. The example implementations herein thus can provide enhanced migration of software services, such as virtual machines or applications, from one host to another with reduced downtime or interruption.

(6) In one example implementation, a method includes receiving ingress traffic transferred by a client node for delivery to an ingress network address associated with an ingress traffic node to a software service, and translating the ingress network address to a provider network address for an initial instance of the software service and transferring the ingress traffic for delivery to the initial instance of the software service with the provider network address. Responsive to a migration trigger event, the method includes withholding transfer of the ingress traffic and buffering the ingress traffic. The method includes obtaining a network address configuration employed for translating among the ingress network address and an updated provider network address for a migrated instance of the software service. Responsive to a migration completion indication, the method includes modifying buffered ingress traffic with the updated provider network address and transferring the buffered ingress traffic for delivery to the migrated instance of the software service.

(7) In another example implementation, a method includes receiving initial egress traffic transferred by an initial instance of a software service for delivery to a destination node through a first egress traffic node selected based at least on a traffic latency between the software service and the first egress traffic node. The method includes receiving initial return traffic transferred by the destination node for delivery to the initial instance of the software service through a second egress traffic node selected based at least on a traffic latency between the software service and the second egress traffic node. Responsive to a migration trigger event, the method includes selecting a third egress traffic node to receive subsequent egress traffic transferred by the migrated instance of the software service for delivery to the destination node, wherein the selection of the third egress traffic node is based at least on a traffic latency between the software service and the third egress traffic node.

(8) In yet another example implementation, an apparatus includes a processing system operatively coupled with one or more computer readable storage media, and program instructions stored on the one or more computer readable storage media. Based at least on being read and executed by the processing system, the program instructions direct the processing system to at least receive ingress traffic transferred by a client node for delivery to an ingress network address associated with an ingress traffic node to a software service, and translate the ingress network address to a provider network address for an initial instance of the software service and transferring the ingress traffic for delivery to the initial instance of the software service with the provider network address. Responsive to a migration trigger event, the program instructions withhold transfer of the ingress traffic and buffer the ingress traffic, and obtain a network address configuration employed for translating among the ingress network address and an updated provider network address for a migrated instance of the software service. Responsive to a migration completion indication, the program instructions modify buffered ingress traffic with the updated provider network address and transferring the buffered ingress traffic for delivery to the migrated instance of the software service.

(9) This Overview is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. It may be understood that this Overview is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) Many aspects of the disclosure can be better understood with reference to the following drawings. While several implementations are described in connection with these drawings, the disclosure is not limited to the implementations disclosed herein. On the contrary, the intent is to cover all alternatives, modifications, and equivalents.
- (2) FIG. 1 illustrates an example system with an application/service migration in an implementation.
- (3) FIG. 2 illustrates example operations for application/service migration in an implementation.
- (4) FIG. 3 illustrates an example system with migration of data/state of an application/service in an implementation.
- (5) FIG. 4 illustrates example operations for data/state migration elements of an application/service in an implementation.
- (6) FIG. 5 illustrates example migration of ingress traffic handling in an implementation.
- (7) FIG. 6 illustrates example operations for migration of ingress traffic handling in an implementation.
- (8) FIG. 7 illustrates example migration of egress traffic handling in an implementation.
- (9) FIG. 8 illustrates example operations for migration of egress traffic handling in an implementation.
- (10) FIG. 9 illustrates an example node system capable of implementing any of the various nodes, control elements, routers, translators, virtualized systems, and associated elements discussed herein.

DETAILED DESCRIPTION

- (11) Migration of virtualized software services can involve moving remote resources, such as virtualized software components, from one host to another. When migration is performed on actively used resources, migration can include targeted goals of reducing service disruptions and downtime to served nodes. Provided herein are various enhancements for reduced downtime migration of virtualized software services across different instantiations, which may include real-time migration over different cloud/server providers and platforms, physical locations, network

locations, and across different network routing elements. The example implementations herein thus can provide enhanced migration of software services (such as virtual machines, virtualized containers, or virtualized applications) from one host to another with reduced downtime and service interruptions.

(12) Various techniques have been developed for accessing and synchronization of remote resources over network links, such as third-party databases, file synchronization services, custom application programming interfaces (APIs), or bespoke synchronization protocols. Current solutions for accessing, synchronization, and migration of resources over a network are characterized by application-specific protocols and interfaces, which result in fragmentation and barriers to adoption. A common technique is using an API in the Linux module `userfaultfd`, a user-space mechanism for handling paging and memory tracking. The examples herein can address problems and shortcomings of existing techniques and use of `userfaultfd`, in part, by presenting a universal approach that enables direct operation on a memory region, circumventing the need for custom-built solutions. This can provide for a solution that is suitable for both local area network (LAN) and wide area network (WAN) environments, among others, by using an approach based on block devices in user space with background push and pull mechanisms. The examples herein can provide a unified API that enables mounting and migration of nearly any state over a network accompanied by manageable changes to existing applications.

(13) Some examples herein can employ the open-source `r3map` (remote mmap) library, but the examples are not limited to such libraries. Particularly, `mmap` is a Linux system call, used for mapping files or devices into memory, enabling a variety of tasks like shared memory, file I/O, and fine-grained memory allocation. `Mmap` can create an operational arrangement with a direct memory mapping between a file and a region of memory. This arrangement means that read operations performed on the mapped memory region directly correspond to reading the file and vice versa, enhancing efficiency as the amount of expensive context switches can be reduced (i.e., more efficient read or write system calls). A significant advantage of `mmap` is its ability to do zero-copy operations. In practical terms, this means that data can be accessed directly as if it were positioned in memory, eliminating the need to copy it from a data storage device (e.g., disk) first. This direct memory access saves time and reduces processing requirements, offering substantial performance improvements. The enhanced migration techniques discussed herein can be categorized into two main types or phases, namely pre-copy migration and post-copy migration.

(14) Pre-copy migration can correspond to a “run-while-copy” nature, meaning that copying of data from an initial host (source host) to a destination host occurs concurrently while the target software service continues to operate. This is also applicable in a generic migration context where other software service state is being updated. In the case of a virtualized software service, such as a VM, the pre-copy migration procedure can start with transferring an initial state of a memory of a VM from an initial host to a destination host. During this operation, if modifications occur to any chunks of data, these chunks are flagged as dirty. These dirty chunks of data are then transferred to the destination until only a threshold amount remains, namely an amount small enough to stay within the allowed maximum downtime criteria corresponding to a transfer time of the remaining chunks from the initial host to the destination host. After this, the VM is suspended at the initial host, enabling the synchronization of the remaining chunks of data to the destination without having to continue tracking dirty chunks. Once this synchronization process is completed, the VM is resumed at the destination host. This pre-copy migration process is discussed herein is reliable, especially in instances where there might be network disruption during synchronization. For instance, at any given point during migration, the VM is available in full either at the initial host or the destination host. A limitation to this approach however is that, if the VM or application changes too many chunks on the initial host during migration, the process might not meet the maximum acceptable downtime criteria. Maximum acceptable downtime is also inherently restricted by the available round-trip time (RTT) between the initial host and the destination host.

(15) Post-copy migration is an alternative live migration approach. While pre-copy migration operates by copying data before halt of a VM, post-copy migration immediately suspends the VM operation on the initial host and resumes it on the destination host, with only a minimal subset of the VM data. During this resumed operation on the destination host, whenever the VM attempts to access a chunk of data not initially transferred during the move, a page fault arises. A page fault, in this context, is a type of interrupt generated when the VM tries to read or write a chunk that is not currently present on the destination host. This triggers the VM system to retrieve missing chunks from the initial host, enabling the VM to continue its operations. One advantage of post-copy migration is that it eliminates the necessity of re-transmitting chunks of “dirty” or changed data before hitting a maximum tolerable downtime. This process can thus decrease this factor and also reduce the amount of network traffic between initial host and destination host. Despite this benefit, post-copy migration could also potentially lead to extended migration times as a consequence of its fetch-on-demand model for retrieving chunks. This is because fetch-on-demand can be highly sensitive to network latency and round-trip time (RTT). Unlike the pre-copy model, this also means that the VM is not available in full on either the initial host or the destination host during migration, requiring potential recovery solutions if network connectivity is lost during the process.

(16) Turning now to the Figures which illustrate various aspects of example implementations, FIG. 1 is presented. FIG. 1 illustrates system **100** comprising one or more client nodes **110** which can communicate with one or more host nodes **120** and **140** over associated network links **191-193**, and network traffic node(s) **190**. Various networks, local area networks (LANs), wide area networks (WANs), and other arrangements can be included with any among links **191-193**. One or more control nodes or command elements might be included to initialize software services and initiate migration operations, and portions of these control nodes or command elements can included in client nodes, host nodes, or separate nodes. Although separate control nodes are omitted from FIG. 1, further Figures illustrate examples of such elements.

(17) Host node **120** executes an initial instance or source instance of a software application or software service, illustrated as software service **135** within isolated environment **131** of virtual machine (VM) host **130**. Host node **140** executes a migrated instance or destination instance of the software application or software service, illustrated as software service **155** within isolated environment **151** of virtual machine (VM) host **150**. Software applications or software services, such as software services **135/155** and other applications and services discussed herein, are referred to herein as “software services” for clarity. However, this terminology does not limit the type of software or application represented in the Figures.

(18) In operation, client node **110** can communicate with an initial instance of software service **135** executed at host node **120**. A migration process can occur which migrates software service **135** from an initial instance at host node **120** to software service **155** as a subsequent migrated instance at host node **140**. As such, software service **155** can thus be conceptualized as a continuation of the initial instance of software service **135** at a different location, server, hardware set, execution environment, or other arrangement. Various data and state information is migrated, such that software service **135** can resume execution as software service **155** without interruption or substantial interruption to client node **110**. This migration process can be referred to as a live migration, shown by operation **101** in FIG. 1. A “live” or “real-time” migration occurs as software service **135** continues to execute up until a migration event, and then resumes the execution state after migration completes—all the while client node **110** is being served by software service **135/155**. Non-live migrations can also be applicable to the examples herein, such as when a client-service connection is halted or broken before migration, and a new connection or service is then instantiated after migration.

(19) To perform migration operation **101**, FIG. 2 is presented listing example operations **200**. Optionally, point ‘A’ in FIG. 2 can connect operationally to other Figures, such as FIGS. 4, 6, and 8. In operation **201**, host node **120** instantiates an initial instance of a software service. As shown in

FIG. 1, software service **135** is initially executed at host node **120** within isolated environment **131** as a part of a virtualized environment, e.g., VM host **130**. As a part of the instantiation, various emulated or virtualized data storage and handling elements can be included, such as a data storage device and a main memory device to support a virtualized processor device that executes software service **135**. Data storage element **133** (“data”) can act as a disk drive, solid state drive, or other mass data storage device. Main memory element **132** (“mem.”) can act as a main memory device or random-access memory (RAM) for a virtualized processor for VM host **130**. As software service **135** executes and otherwise operates (operation **202**), contents of both memory element **132** and data element **133** are updated and/or changed. This operation of software service **135** can include providing services to client node **110**, such as web services, data services, processing services, media content services, or other services which may include display of various content and user interaction via user interface **111** of client node **110**.

(20) At some point during the operation of software service **135**, a migration might be desired to move the initial instance of software service **135** to another hardware device, server system, geographical location, cloud data services provider, data center, or other different physical or logical location. Responsive to this intent to migrate, pre-migration activities might be triggered, or these pre-migration activities might be performed on an ongoing basis once software service **135** is initialized. In operation **203**, these pre-migration activities can be performed by VM host **130** or a separate control entity (not shown, see examples for node **520**, node **770**, and system **1000** below), which can include pre-migrating data and various state (such as execution state or virtual environment state) associated with software service **135**, as well as identifying possible locations for the migration to land.

(21) Among the pre-migration activities, data stored by memory element **132** and data element **133** can be synchronized to a storage location separate from host **120**, such as over a network connection to a remotely located data storage system. This synchronization can occur periodically or responsive to changes made to the contents of memory element **132** and data element **133**, including combinations thereof. Thus, pre-migration data images might be stored off-site from host **120** to be ready for any future migration activities.

(22) To enable these synchronization activities, such as for memory element **132** which typically contains active execution data and pages for software service **135** (as opposed to long-term data storage or paged-out data in data element **133**), various techniques can be applied. One example includes pull-based data synchronization with UNIX/Linux userfaultfd functions in a scenario where pre-migration is not employed. Thus, the userfaultfd functions allow for the implementation of a post-copy migration scenario. In this setup, a data storage region is created on a remote storage system. A migrated software service can then start to read from this remote region after being resumed and page faults can be triggered by normal operation, which are resolved by fetching the relevant data offset from the remote region. However, userfaultfd has drawbacks which limit its functionality for pre-migration of software service **135**. Instead, other enhanced examples discussed herein can employ UNIX/Linux mmap functions to provide pre-migration and other migration activities.

(23) Advantageously, mmap functions can provide for pre-migration scenarios and push-based synchronization. Mmap allows mapping a main memory region to a file, such that main memory (e.g., memory element **132**) is implemented by this file. Mmap can be used can be used to synchronize data/state of an actively executed software service (e.g., an application) to a data storage region created on a storage system comprising the file. Since the main memory of software service **135** is mapped to a file, and likewise the storage region stores the file, when writes happen to the storage region. Writes are detected and copied to a remote data storage system, establishing at least a portion of a pre-copy (pre-migration) scenario. In some examples, writes done to a mmap linked storage region are not immediately written to the underlying file, since the kernel still might employ caching on a mmap-ed region in order to speed up reads/writes. In such examples, an

msync syscall can be used, which works similarly to the sync syscall by flushing any remaining changes from the cache to the underlying file. Instead of using inotify or a similar event-based system to track changes, a polling system can be used. This has drawbacks, namely latency and computational load.

(24) Another example using an mmap-based approach for both pre- and post-copy migration is to mmap a block device instead of a file. This block device can be provided through a variety of APIs, most notably a Linux network block device (NBD). By providing a NBD device through a kernel NBD client, the NBD device can be connected to a remote NBD server, which hosts the resource as a memory region. Reads/writes from/to the mmap-ed memory region are resolved by the NBD device, which forwards it to the client, which then resolves them using the remote server. This approach does not actually copy the memory region to the destination host, but rather comprises a mount of a remote memory region over the NBD protocol. A benefit of using mmap on such a block device (NBD) instead of a file on a custom file system (as in the previous example) is reduced complexity. For the use case of memory synchronization, not all the features provided by a full file system are required, which means that the implementation of an NBD server and client, as well as the accompanying protocols, is significantly less complex and can also reduce the overhead of the system.

(25) Thus, through one of the various techniques mentioned above, a pre-migration process ensures that live execution data and/or application state for software service **135** is synchronized to a file or block device, which can then further be synchronized to a remote storage system over a network link or other remote link. At some point during operation of software service **135**, a migration operation is triggered (operation **204**). The migration can be triggered automatically due to various criteria or manually with operator intervention, and signaled to VM host **130** over a network link, API, or other remote signaling. Responsive to the migration trigger, VM host **150** can be instantiated at host **140**, which can then establish isolated environment **151**. Software service **155** can be instantiated within isolated environment **151**, and data/state corresponding to software service **135** can be employed for software service **155** to resume operations of software service **135**. In this manner, software service **155** comprises a migrated instance of software service **135**. In addition to data/state migrations, network properties and connections can be migrated, and this network migration will be discussed in more detail in later Figures.

(26) With regard to data and state migration, operations **205** and **206** relate to post-copy migration of data and state (associated with execution of software service **135**) to a migrated instance comprising software service **155**. This includes halting or stopping the initial instance of the software service to prevent further changes to data/state and copying an incremental update to the data/state not already synchronized during pre-migration activities. These incremental updates are typically small in size, such that the time to copy or synchronize to a migrated instance of the software service is on the order of hundreds of milliseconds (ms). Once the incremental transfers occur, then the migrated instance of the software service (e.g., software service **155**) can resume operation using the data/state from both the incremental updates (post-copy migration) and pre-migrated data (pre-copy migration). Once initiated at host **140**, software service **155** can continue operation of software service **135** using the migrated data/state and network properties. The two-prong pre/post migration approach ensures a fast transition among instances of the software service and provides uninterrupted service to client node **110**.

(27) In addition to the data/state migration, operation **207** includes migration of network state, connections, or properties from the initial instance of the software service (e.g., software service **135**) to the migrated instance of the software service (e.g., software service **155**). This migration of network state includes informing network routing or handling elements to reference new network addressing properties for software service **155** which are updates to network addressing properties from software service **135**. Specifically, TCP/IP addressing for ingress and/or egress traffic nodes that handle inbound/outbound traffic between client node **110**, software service **135**, and various

network interfacing elements of host **120**, VM host **130**, and isolated environment **131** is updated such that client node **110** maintains the same network addressing used to access the software service across both the initial instance and the migrated instance. Thus, client node **110** need not be aware of the migration process and can receive uninterrupted service with respect to the software service without regard to if hosted at host **120** or host **140**.

(28) Returning now to a discussion on the elements of FIG. **1**, client node(s) **110** comprise various computing nodes having user interface elements, such as U/I **111**, network interface controllers and network interfacing elements, data processing and storage elements, and other various components. Examples include computers, servers, tablet computing devices, media streaming devices, smartphone devices, laptops, desktop computers, gaming consoles, end user devices, customer equipment, mobile Internet appliance, media player, endpoint terminals, kiosk nodes, or other various computing devices, including combinations thereof. In some examples, client node(s) **110** execute software, such as operating systems, display drivers, and applications to display graphical content on a display screen. Client node(s) **110** can include one or more CPUs or microprocessors and other processing circuitry that retrieves and executes software, and any number of end user applications, from an associated storage system. Each processing element can be implemented within a single processing device but can also be distributed across multiple processing devices or sub-systems that cooperate in executing program instructions. Examples of each processing module include general purpose central processing units, application specific processors, and logic devices, as well as any other type of processing device, combinations, or variations thereof.

(29) Network traffic node(s) **190** comprise one or more network routing, handling, traffic processing, and transport elements or equipment. Network traffic node(s) **190** can comprise ingress traffic nodes, egress traffic nodes, network routers/switches, network traffic caching nodes, and other equipment. Network traffic node(s) **190** can include various network switching circuitry to communicatively link individual network links to other network links based on network addressing, traffic patterns, network status, or other traffic properties. In one example, network traffic node(s) comprise Transmission Control Protocol/Internet Protocol (TCP/IP) routing and control elements, which can include Ethernet switching and routing elements corresponding to wired or wireless connections, which can refer to any of the various network communication protocol standards and bandwidths available, such as 10BASE-T, 100BASE-TX, 1000BASE-T, 10GBASE-T (10 GB Ethernet), 40GBASE-T (40 GB Ethernet), gigabit (GbE), terabit (TbE), 200 GbE, 400 GbE, 800 GbE, Fifth Generation New Radio (5G NR), Long Term Evolution (LTE), Institute of Electrical and Electronic Engineers (IEEE) 802.11 (Wi-Fi), or other various wired and wireless formats and speeds.

(30) Hosts **120** and **140** each comprise computing platforms which can host various software payloads for execution, such as applications, software services, user/client payloads, software workloads, virtual machines, virtualization environments, containers, or other software elements. Hosts **120** and **140** comprise hardware elements including processing elements, data storage and memory elements, network interfaces, and optional user interface elements. Examples of hosts **120** and **140** include servers, blade servers, computers, elements of a distributed computing system, or other computing systems. Typically, hosts **120** and **140** will each include a motherboard or other system circuit board having a central processing unit (CPU) coupled thereto along with memory devices, such as random-access memory (RAM) or dynamic RAM (DRAM). The CPU can be a component in a processing system formed from among one or more microprocessor elements. Various peripheral devices can physically couple to hosts **120** and **140** via corresponding links within an enclosure or chassis. These peripheral devices can include graphics cards housing graphics processing units (GPUs), data storage drives using various computer-readable media, network interface controllers (NICs) having physical layer elements to couple to network links (e.g., Ethernet), or other devices including user interface devices.

(31) Hosts **120** and **140** also include various software which is executed by one or more physical

processors and operating systems. Software typically includes host operating systems, user applications, device drivers, user data, hypervisor software, telemetry software, or other various software elements. The hardware and host software elements (e.g., host operating system) can be isolated or virtualized from payload software elements by a virtualization system, which can include one or more virtual machines (VMs), virtualized environments, or containers, among other virtualized arrangements, which abstract the physical hardware and host software elements from the payload/virtualized software. Among the elements which provide such functionality include hypervisors or other software elements that virtualize and apportion hardware resources among the payloads presently executing. As noted, these resources can include CPUs, GPUs, network interfaces, RAM, mass storage drives (e.g., hard disk drives (HDDs) or solid-state drives (SSDs)), or other hardware elements.

(32) Although not required, hosts **120** and **140** are also each associated with a different provider in this example. Specifically, host **120** is associated with provider A and host **140** is associated with provider B. Providers can correspond to different service providers, cloud computing providers, hosting companies, service platforms, terms of service regions/agreements, or other arrangements. In some examples, providers can correspond to the same entity but having different sets of physical hardware dedicated to different service levels or qualities of service, such as different tiers of service for different pricing, bandwidth limits, processing limits, or other criteria. Thus, migrating a software service from a first provider to a second provider might change physical locations within a hosting environment, logical locations within a virtualized environment, geographic locations to different server facilities, higher/lower terms of service or quality of service limits, different network locations with different network latencies, locations, or bandwidths, from stationary to mobile/moving arrangements, from home networks to roaming networks, or other changes which prompt or trigger a migration event.

(33) VM host **130** and **150** comprise virtualized environments each including at least one among a virtual machine, operating system level virtualization, application-level virtualization, system virtual machine, process virtual machine, container, or containerized environment which executes on hardware provided by host nodes **120** and **140**. VM host **130** and **150** can instantiate any number of instances of software services, such as software services **135** and **155**. These software services might comprise further virtualized environments nested within the virtualized environment of VM host **130** and **150** (e.g., nested virtual machine, operating system level virtualization, application-level virtualization, system virtual machine, process virtual machine, container, or containerized environment). Isolated environments **131** and **151** can isolate a network namespace within VM host **130** and **150** for software services **135** and **155**. This can include isolating more than just network properties, and can be configured to isolate software services **135** and **155** from differences in hardware clock speeds or processor speeds and other hardware aspects from virtualized instance of the software services, provide consistent network addressing for virtualized instances of software services across various machines/locations and migrations.

(34) Also, interconnection links and communications carried by such links in FIG. 1 (and other Figures herein), can be encrypted. This encryption can take various forms, including data encryption of data payloads, packet encryption, encapsulation of packets and encryption thereof, encryption of links and traffic carried thereon, and other encryption techniques. These techniques can include end-to-end encryption, various authentication and key exchange technologies, asymmetric encryption, symmetric encryption, AES (Advanced Encryption Standard) encryption types, Data Encryption Standard (DES) types, hashing, public-key encryption, private-key encryption, and various other types or techniques of encrypting data, packets, links, interfaces, network stacks, and the like.

(35) Turning now to FIG. 3, a further system illustration is provided detailing data migration and state migration for a software service, such as a virtual machine, container, or other virtualized software element which may comprise an application, software service, software server element, or

other software element **FIG. 3** illustrates system **300** comprising one or more client nodes **310** which can communicate with one or more host nodes **320** and **340** over associated network links **391-393**, and network traffic node(s) **390**. Various networks, local area networks (LANs), wide area networks (WANs), and other arrangements can be included with any among links **391-393**. One or more control nodes or command elements might be included to initialize software services and initiate migration operations, and portions of these control nodes or command elements can be included in client nodes, host nodes, or separate nodes. Although separate control nodes are omitted from **FIG. 3**, further Figures illustrate examples of such elements.

(36) Also included in **FIG. 3** is data storage system **360**. Data storage system **360** can be configured to store pre-migrated data/state for software services, and provide stored data/state to migrated instances. Data storage system **360** includes data storage unit **361** which can store both data and state associated with instances of software services, such as corresponding to main memory contents and mass storage device contents of memory element **332** and data storage element **333** of the initial instance of software service **335**. Data storage system **360** can communicate with network interface elements of hosts **320** and **340** over one or more network links (not shown). Although shown as a remote or off-site data storage system in **FIG. 3**, portions of data storage system **360** might instead be included among hosts **320** or **340**, or other locations.

(37) Host node **320** executes an initial instance or source instance of a software application or software service, illustrated as software service **335** within isolated environment **331** of virtual machine (VM) host **330**. Host node **340** executes a migrated instance or destination instance of the software application or software service, illustrated as software service **355** within isolated environment **351** of virtual machine (VM) host **350**. Host **320** includes hypervisor (HV) **325** and host **340** includes HV **345**. These hypervisors provide environments for virtual machines (VMs), containers, or other virtualized elements to operate and share underlying physical computing resources, such as CPUs, GPUs, RAM/memory, network interfaces, and mass storage devices (SSDs/HDDs) of a corresponding server. HVs **325** and **345** are representative of a virtualized execution system, which includes virtualized user spaces having individual operating systems and applications, such as for software services **335/355**. As seen in **FIG. 3**, host **320** includes physical hardware comprising CPU **321**, main memory (RAM) **322**, data storage device (data) **323**, and network interface controller (NW) **324**, among other elements. Likewise, host **340** includes physical hardware comprising CPU **341**, main memory (RAM) **342**, data storage device (data) **343**, and network interface controller (NW) **344**, among other elements.

(38) In operation, client node **310** can communicate with an initial instance of software service **335** executed at host node **320**. A migration process can occur which migrates software service **335** from an initial instance at host node **320** to software service **355** as a subsequent migrated instance at host node **340**. As such, software service **355** can thus be conceptualized as a continuation of the initial instance of software service **335** at a different location, server, hardware set, execution environment, or other arrangement. Various data and state information is migrated, such that software service **335** can resume execution as software service **355** without interruption or substantial interruption to client node **310**. Software service **335** can continue to execute up until a migration event, and then resume the execution state after migration completes—all the while client node **310** is being served by one among software services **335/355**.

(39) To perform a migration operation, **FIG. 3** includes example copy or migration operations with respect to data storage system **360**. These operations are further described in example operations **400** of **FIG. 4**. In operation **401**, host node **320** instantiates an initial instance of a software service. As shown in **FIG. 3**, software service **335** is initially executed at host node **320** within isolated environment **331** as a part of a virtualized environment, e.g., VM host **330**. As a part of the instantiation, various emulated or virtualized data storage and handling elements can be included, such as a data storage device and a main memory device to support a virtualized processor device that executes software service **335**. Data storage element **333** (“data”) can act as a disk drive, solid

state drive, or other mass data storage device. Main memory element **332** (“mem.”) can act as a main memory device or random-access memory (RAM) for a virtualized processor for VM host **330**. As software service **335** executes and otherwise operates, contents of both memory element **332** and data element **333** are updated and/or changed. This operation of software service **335** can include providing services to client node **310**, such as web services, data services, processing services, media content services, or other services.

(40) At some point during the operation of software service **335**, a migration might be desired to move the initial instance of software service **335** to another hardware device, server system, geographical location, cloud data services provider, data center, or other different physical or logical location. Responsive to this intent to migrate, pre-migration activities might be triggered, or these pre-migration activities might be performed on an ongoing basis once software service **335** is initialized. In operation **402**, these pre-migration activities can be performed by VM host **330** or a separate control entity (not shown), which can include pre-migrating data and various state (such as execution state or virtual environment state) associated with software service **335**, as well as identifying possible locations for the migration to land.

(41) Among the pre-migration activities, data stored by memory element **332** and data element **333** can be synchronized to a storage location separate from host **320**, such as over a network connection to data storage system **360**. This synchronization can occur periodically or responsive to changes made to the contents of memory element **332** and data element **333**, including combinations thereof. To enable these synchronization activities, such as for memory element **332** which typically contains active execution data and pages for software service **335** (as opposed to long-term data storage or paged-out data in data element **333**), various techniques can be applied.

(42) FIG. 3 employs UNIX/Linux mmap and block device functions to provide pre-migration and other migration activities. Mmap functions can provide pre-migration and push-based synchronization. Mmap allows mapping a main memory region to a file, such that main memory (e.g., memory element **332**) is implemented by this file. Mmap can be used to synchronize data/state of an actively executed software service (e.g., an application) to a data storage region created on a storage system comprising the file. The main memory of software service **335** is mapped to a file, as shown for label “1: mapshare” in FIG. 3. Storage region **326** of data storage device **323** stores the file. In some examples, writes done to a mmap linked storage region are not immediately written to the underlying file, since the kernel still might employ caching on a mmap-ed region in order to speed up reads/writes. In such examples, an msync syscall can be used, which works similarly to the sync syscall by flushing any remaining changes from the cache to the underlying file. Instead of using inotify or a similar event-based system to track changes, a polling system can be used. This has drawbacks, namely latency and computational load.

(43) In addition to using an mmap-based approach for both pre- and post-copy migration, a block device can be employed to host the mmap file. This block device can be provided through a variety of APIs, most notably a Linux network block device (NBD). By providing an NBD device through a kernel NBD client, the NBD device can be connected to a remote NBD server, which hosts the resource as a memory region. In FIG. 3, storage region **326** of data storage device **323** comprises an NBD and includes at least a file implementing or synchronous with memory element **332**.

Writes to block device **326** are detected and copied to a remote data storage system, namely data storage system **360** in view **301**, which is also shown by label “2: cyclic updates (pre-copy)” in FIG. 3. Data storage system **360** hosts an NBD in storage region **366** in data storage device **361**, establishing at least a portion of a pre-copy (pre-migration) copy of block device **326**. As shown, block device **326** can house data for both memory element **332** and data storage element **333**, although some examples might include more than one block device. Thus, a pre-migration process ensures that live execution data and/or application state for software service **335** is synchronized to a file or block device, which can then further be synchronized to a remote storage system (**360**) over a network link or other remote link.

(44) At some point during operation of software service **335**, a migration operation is triggered (operation **403**). The migration can be triggered automatically due to various criteria or manually with operator intervention, and signaled to VM host **330** over a network link, API, or other remote signaling. Responsive to the migration trigger, VM host **350** can be instantiated (operation **404**) at host **340**, which can then establish isolated environment **351**. Software service **355** can be instantiated within isolated environment **351**, and data/state corresponding to software service **335** can be employed for software service **355** to resume operations of software service **335**. In this manner, software service **355** comprises a migrated instance of software service **335**. In addition to data/state migrations, network properties and connections can be migrated, and this network migration will be discussed in more detail in later Figures.

(45) In FIG. **3**, once software service **355** is instantiated in isolated environment **351** of VM host **350**, then a real-time data migration of pre-copy and post-copy data/state can occur (operation **405**). For pre-copy data/state (operation **406**), a link to NBD in storage region **366** can be established for software service **355**. Label “3: dynamic link” in FIG. **3** shows this initial linking or mapping of the NBD in storage region **366** to an NBD in storage region **346** of data storage device **343** of host **340**. In this manner, data can be read or served ‘live’ by software service **355** from the NBD in storage region **366** over a network link until a migration has completed. In some examples, client requests can be served by data in NBD in storage region **366** through software service **355**. For post-copy data/state (operation **407**), a final record of the data/state of software service **335** is copied to the NBD in storage region **346** in label “4: state copied (post-copy)” in FIG. **3**. This post-copy data includes previously un-migrated contents of the NBD in storage region **326**. Once all data/state has been transferred, mapped, or linked for use by software service **355** with respect to the NBD in storage region **366**, then data/state of software service **335** can be considered migrated to software service **355**.

(46) Further terminology can be employed for the various elements of FIG. **3**. For example, the NBD in storage region **326** might be referred to as a seeder, and the NBD in storage region **366** referred to as a leecher, or other terminology. Connections can be established among the seeder and leecher, and changes can be tracked at the seeder which are echoed or synchronized unidirectionally to the leecher, using either a push or pull mechanism. For example, background ‘pulls’ might be employed to echo data in a delta- or de-duplicated manner from the seeder to the leecher. Upon migration, NBD in storage region **326** might be suspended from further changes, and a collection of unsynchronized changes (incrementally updated data chunks) or un-echoed changes can be tabulated and transferred to the migrated instance. These unsynchronized changes are typically small in size, such that the time to copy or synchronize to the new/migrated instance of the software service is on the order of ~200 ms. Once the incremental transfers occur, then the migrated instance of the software service (e.g., software service **355**) can resume operation using the data/state from both the incremental updates (post-copy migration) and pre-migrated data (pre-copy migration). Once initiated at host **340**, software service **355** can continue operation of software service **335** using the migrated data/state and network properties. The two-prong pre/post migration approach ensures a fast transition among instances of the software service and provides uninterrupted service to client node **310**.

(47) Thus, FIGS. **1-4** discuss various examples of migrating an execution workload, such as a software service, from one virtualized environment to another virtualized environment with reduced downtime or reduced interruption for client nodes. Among these examples includes implementations that establish a first instance of a software service in a first network namespace of a first virtualized environment configured to use block devices to emulate main memory and a data storage device for the software service. These examples can include periodically synchronizing contents of the block devices to one or more files to reflect the main memory and the data storage device and pre-migrating the one or more files to a migration storage location. Responsive to a migration trigger event, the examples can initiate a migration operation to establish a second

instance of the software service in a second network namespace of a second virtualized environment. The migration operation can comprise directing the second instance of the software service to reference the migration storage location to resume a state of the first instance, transferring un-migrated contents of the block devices for the first instance to block devices of the second instance, and migrating network addressing properties of the first instance such that a client node retains destination network address associated with the first instance to communicate with the second instance.

(48) The examples herein can produce a migration with reduced or minimized downtime, and this downtime can be further optimized by selecting an order in which data is migrated between nodes. For example, writes can be tracked in real-time (without affecting the write performance of the workload), and migration can be provided for the least-volatile blocks of memory first (and most volatile last), such that the amount of dirty data re-transferred is reduced or minimized. The data volatility can be tracked by recency of access, with thresholds established for recency of access/use/write/read to differentiate between most volatile and least volatile. Other metrics, thresholds, and data categorization can be employed. The data volatility can be assessed continuously or periodically, as well as before and during the migration. This can result in higher data transfer overall, but significantly reduces the overall time to complete a migration of the data after the corresponding VM is paused. Specifically, since only the ‘most’ volatile data needs to be transferred after pausing the VM, then the ‘least’ volatile data has already been transferred ahead of time.

(49) Also, in the examples herein, a VM can be paused and migrated without moving the most volatile blocks of data. Those “most volatile” blocks can be left at the source host. The VM is then resumed on the destination host without those blocks being available locally. However, when the destination host attempts to read or write to those volatile blocks, those operations will succeed because the blocks are transferred just-in-time. This can lead to slightly higher latency for the first read at the destination host because some blocks/data are transferred just-in-time (future reads are local to the destination host). This can lead to slightly higher latency as well as for partial block writes (where part of a block is written), due to the need to fetch the full block first and then overwrite the partial areas locally at the destination host (future writes to the block do not have additional latency). The end result is that a VM does not need to be paused for much time-because the operations herein are not sending any disk or memory data, just state and some metadata which is comparatively small. The amount of time the VM is paused is the definition of downtime, so by keeping that close to 0, the examples herein are actively providing an enhanced reduced downtime migration.

(50) In addition to the data/state migration noted above, migration of network state can be performed for network connections or network properties from the initial instance of the software service (e.g., software service 335) to the migrated instance of the software service (e.g., software service 355). FIGS. 5 and 6 provide examples of ingress traffic routing and network migration, where ingress traffic relates to traffic transferred by a client node for delivery to a corresponding software service (and associated return traffic). FIGS. 7 and 8 provide examples of egress traffic routing and network migration, where egress traffic relates to traffic transferred by a software service for delivery to a third party, client node, or other application or service (and associated return traffic).

(51) Turning now to FIG. 5, migration of network connections and network properties for a software service is shown in one example implementation. Specifically, FIG. 5 illustrates migration of ingress traffic routing for a software service migrating from a first instance to a second instance. Ingress traffic refers to network traffic transferred for delivery to a software service. This ingress traffic might be originated by a client node or other nodes that interface with a software service. To transfer such traffic, an originating node typically includes network addressing among packets or frames to reach the software service. This network addressing, as will be discussed below, is

abstracted from an actual network address of a software service, which typically is instantiated within an isolated environment, such as a network namespace within a virtualized environment or virtual machine. By abstracting the network addressing for a software service, a client node can continue to communicate using the same network address to reach a software service both before and after migration of the software service from one instance to another instance.

(52) One example of ingress traffic handling is shown in system **500** of FIG. 5. System **500** includes client node **510**, control node **520**, ingress traffic nodes **521-523**, VM host **530**, and VM host **560**. Client node **510** can communicate with any among ingress traffic nodes **521-523** over network link **580**. Control node **520** can communicate with any among ingress traffic nodes **521-523** over network link **581**, shown in a simplified representation in FIG. 5. Ingress traffic nodes **521-523** can communicate with various nodes, such as with VM hosts **530** and **560** over associated network links, such as network link **582** shown for ingress traffic node **522**. VM hosts **530** and **560** can communicate over network links **583** and **584** coupled to router elements **531** and **561**. Router elements **531** and **561** can be virtualized network router elements corresponding to physical network interfaces (not shown) of host computing devices or host compute units.

(53) Client node **510** comprises various computing nodes having user interface elements, network interfacing elements, data processing and storage elements, and other various components. Examples include those discussed above for client node(s) **110** above, such as various types of computing devices or media playback devices, although variations are possible.

(54) Ingress traffic nodes **521-523** comprise one or more network routing, handling, traffic processing, and transport elements or equipment. Ingress traffic nodes **521-523** can include various network switching circuitry to communicatively link individual network links to other network links based on network addressing, traffic patterns, network status, or other traffic properties. In this example, ingress traffic nodes **521-523** include translator elements **524-526** and buffer elements **527-529**. Translator elements **524-526** comprise network route translation elements, including software components, data processing elements, routing table storage elements, network router elements, and other components. Buffer elements **527-529** comprise data storage elements, such as memory elements, mass storage elements, SSDs, HDDs, RAM, or other data storage components capable of storing network traffic in the form of packets, frames, payloads, headers, addressing, and other portions of network traffic.

(55) Control node **520** comprises control elements for instructing ingress traffic nodes **521-523** and VM hosts **530** and **560** with regard to the various operational scenarios described herein. Control node **520** can include one or more microprocessors and other processing circuitry that retrieves and executes software, and any number of control software applications, from an associated storage system. Each processing element can be implemented within a single processing device but can also be distributed across multiple processing devices or sub-systems that cooperate in executing program instructions. Examples of each processing module include general purpose central processing units, application specific processors, and logic devices, as well as any other type of processing device, combinations, or variations thereof. In FIG. 5, control node **520** can take inputs from various external elements and nodes, and perform various control operations based at least on these inputs.

(56) As shown in FIG. 5, software service **545** is initially executed at a host node within isolated environment **540** as a part of a virtualized environment, e.g., VM host **530**. Software service **555** within isolated environment **550** is also included within VM host **530**, which might comprise another parallel instance of software service **545** that services other client nodes or other services/applications for client node **510**. VM host **530** also includes router element **531** which comprises a virtualized router configured to route or otherwise handle network traffic between external network links/nodes and network elements internal to VM host **530**, such as those for isolated environments **540** and **550**. Furthermore, VM hosts **530** and **560** might further include virtualized environments within their own virtualized environments. These embedded, recursive, or

nested virtualized environments can include those to instantiate software services **545** and **555**. Thus, software services **545** and **555** can comprise a further virtualized environment nested within another virtualized environment (e.g., VM hosts **530** and **560**), where the further virtualized environment has a network interface and instantiates a client-facing software service.

(57) VM hosts **530** and **560** can be hosted by physical computing platforms, such as shown for hosts **120** and **140**, although variations are possible. Thus, these physical computing platforms might include various hardware elements including processing elements, data storage and memory elements, network interfaces, and optional user interface elements. The hardware and host software elements (e.g., host operating systems) can be isolated or virtualized from payload software elements by a virtualization system, which can include one or more virtualized environments represented by isolated environments **540** and **550** in FIG. 5.

(58) FIG. 5 shows various functional elements which are interconnected by network links. The actual network links are omitted from FIG. 5 for most of the functional elements, but are shown as exemplary links **580-584**. Links **580-584** can comprise any type of network link formed from one or more physical links comprising electrical, optical, or wireless signaling layers. Various network protocols and traffic types can be carried by such links, such as TCP/IP or other similar protocols associated with various network stacks and transport/internet layers. The examples herein are not intended to limit the links to any particular protocol or network link type. Also, these network links can carry various frames, packets, datagrams, or other granular data representation, which can vary based on the exact link/protocol employed. Typically, addressing used for routing the traffic is included in a header portion of a packet or frame, with payload portions configured to carry data or other packets/frames with associated addressing.

(59) Turning now to various operational scenarios for FIG. 5, operations **600** of FIG. 6 are presented. In operation **601**, an initial instance of a software service is established, such as shown for software service **545** in isolated environment **540** of VM host **530**. Along with instantiating software service **545** and various storage/memory elements discussed above, corresponding network interfaces and network states can be established. These network states include various network addressing, network properties of isolated environment **540**, network routing tables, ingress/egress configurations, and other states and properties. In one example, ingress traffic nodes and router elements of the VM hosts receive network address configurations from control node **520**, where the network address configurations are determined based at least on a provider network address for an initial instance of the software service. A provider or other service entity that hosts VM host **530** can provide network addressing (e.g., IP addressing) for some elements used by VM host **530**, which can correspond to physical network interfaces of a host computer system, those of router element **531**, and those associated with isolated environments **540** and **550**. Other addressing can be associated with such elements, such as media access control (MAC) addressing which typically corresponds to a physical or virtual network interface controller, and may be assigned by other entities than control node **520** or the service provider of VM host **530**.

(60) After network states have been configured/assigned for ingress traffic nodes and software services within any corresponding virtualized environments, then network traffic might be issued by client nodes for receipt by the software services. For example, client node **510** can originate or otherwise transfer network traffic, such as service requests, data requests, streaming media commands, or other traffic, over link **580** for delivery to software service **545** using an assigned network address. Moreover, this network address might be associated with a particular ingress traffic node, such as ingress traffic node **522** based on various factors including proximity (physical or network routing), latency, availability, a type of software service associated with software service **545**, a type of client node associated with client node **510**, network status, congestion, network routing considerations, and other factors. As seen in FIG. 5, ingress traffic node **522** might have a network address of 40.40.40.40 associated therewith, while ingress traffic node **521** has 40.40.40.41 and ingress traffic node **523** has 40.40.40.42. Client node can use 40.40.40.40 to on-

ramp network traffic intended for software service **545**, and as will be discussed below, can maintain use of this network address during and after migration of software service **545** to another location/server.

(61) Thus, ingress traffic node **522** receives (operation **602**) ingress traffic transferred by client node **510** for delivery to an ingress network address (e.g., 40.40.40.40) associated with a selected ingress traffic node (**522**) to software service **545**. Ingress traffic node **522** translates (operation **603**) the ingress network address to a provider network address for an initial instance of software service **545** and transfers the ingress traffic for delivery to the initial instance of the software service **545** with the provider network address. In this example, the provider network address corresponds to a network address assigned to router element **531** of VM host **530**, namely 1.1.1.1. Translator **525** of ingress traffic node **522** can perform this network address translation, using various techniques including NAT (network address translation) of IP addressing from a first address (40.40.40.40) to a second address (1.1.1.1). Translator **525** can maintain various translation tables or other data structures for this purpose. From here, ingress traffic node **522** transfers or forwards (operation **604**) the ingress traffic for delivery to router element **531** over link **582**, which can optionally include altering the source IP address to 40.40.40.50 as shown in FIG. 5 to provide for return traffic routing to ingress traffic node **522** instead of client node **510**.

(62) Router element **531** of VM host **530** can then receive the ingress traffic which has the provider network address over link **583** and determine a further destination for the ingress traffic. These further destinations can include isolated environment **540** or **550** in VM host **530**, and several network address translation elements are encountered along these pathways. These address translations provide network isolation for software services within virtualized environments, abstract the network addressing and interfacing from physical hardware resources, and provide for enhanced migration of the software services to different physical hardware resources. Specifically, router element **531** translates the provider network address into a virtual environment network address associated with an initial isolated namespace instantiating the software service. For software service **545**, the virtual environment network address is shown as 10.0.0.1 which is associated with isolated environment **540** comprising a namespace housing further translation elements and software service **545**. For software service **555**, the virtual environment network address is shown as 10.0.0.3 which is associated with isolated environment **550** comprising a namespace housing further translation elements and software service **555**.

(63) In one example, router element **531** then transfers the ingress traffic having the virtual environment network address, which is received by translator element **541** of isolated environment **540**. Translator element **541** is configured to interface into the initial isolated namespace (isolated environment **540**), and translates the virtual environment network address (10.0.0.1) to an internal namespace network address, shown in FIG. 5 as 10.0.0.2. Translator element **541** then transfers the ingress traffic having the internal namespace network address for receipt by another translator element **542**. In second translator element **542**, another translation occurs to translate the internal namespace network address (10.0.0.2) to a software service network address, shown in FIG. 5 as 172.100.100.1. Translator element **542** then transfers the ingress traffic having the software service network address for receipt by software service **545**.

(64) As noted above, network traffic issued by client node **510** can propagate through a series of network links, connections, routers, and translation elements to reach software service **545**, which may itself be implemented as a virtualized environment, container, or virtual machine. Software service **545** has network addressing associated therewith, such as the software service network address (172.100.100.1) as well as an address used for return traffic of 172.100.100.100, and associated MAC addressing (e.g., 66:66:66:66:66), which can vary based on implementation. However, to provide the enhanced migration of software service **545** to other location, machine, virtual environment, provider, server, or other instance hardware/software platform, the examples herein maintain the same network addressing for software service **545** before, during, and after

migration. Moreover, the same network addressing used by client node **510** to reach software service **545** is maintained as the same network addressing before, during, and after migration. As will be discussed, the combination of the ingress traffic node address translation, as well as the use of multiple translator elements in the VM host entities provides for these enhanced operations. (65) Specifically, a migration trigger is initiated or detected (operation **605**), which triggers various network state migration activities for a software service (as well as data/state migration noted herein). In ingress traffic node **522**, this trigger event prompts ingress traffic node **522** to withhold transfer of additional or new ingress traffic and buffer (operation **606**) any intervening ingress traffic within traffic buffer **528**. Before or during the migration timeframe, execution of software service **545** is paused/halted at VM host **530** and software service **545** is instantiated at VM host **560** within isolated environment **570**. Various virtualized environment state/data or storage/memory data can be migrated (pre- or post-migration, noted above), and network state can be configured to support the new migrated instance at VM host **560**. However, different network addressing will be associated with reaching VM host **560** and migrated software service **545** therein.

(66) Control node **520** can receive this different network addressing associated with VM host **560** from a provider or other service node corresponding to VM host **560**, and proceed to update various other nodes with this network addressing. For instance, ingress traffic node **522** receives an updated network address configuration from control node **520**, where the updated network address configuration is determined based at least on an updated provider network address for the migrated instance of the software service. The updated network address configuration can be employed for translating among the ingress network address provided in traffic from client node **510** and an updated provider network address for a migrated instance of software service **545** at VM host **560**.

(67) Once the updated network configuration is obtained by ingress traffic node **522**, and any corresponding migration completion indication has been received/notified (operation **607**), then ingress traffic node **522** can resume translating and transferring traffic for delivery to software service **545** using the updated network configuration in translator **525**. Ingress traffic node **522** transfers a manifest of the buffered and intervening ingress traffic destined for software service **545**. Ingress traffic node **522** can modify buffered ingress traffic with the updated provider network address and transfer the buffered and intervening ingress traffic for delivery to the migrated instance of software service **545** (operation **608**). Typically, this buffering is of a short duration (e.g., 100-200 ms), and client node **510** does not observe any appreciable delays in traffic reaching software service **545** during the migration. The migration completion indication denotes that the migrated instance of the software service has started at the new location/machine/provider. After this indication, ingress traffic node **522** can receive further (new) ingress traffic transferred by client node **510** for delivery to the same ingress network address used before migration, then translate the ingress network address to the updated provider network address for the migrated instance of the software service, and transfer the further ingress traffic for delivery to the migrated instance of software service **545** with the updated provider network address. Advantageously, in addition to no interruption in service, client node **510** employs the same unchanged ingress network address (40.40.40.40) to communicate with the initial instance of the software service and the migrated instance of the software service. Thus, the migration can be considered unrevealed and transparent to client node **510**.

(68) From here, ingress traffic node **522** transfers or forwards the ingress traffic for delivery to router element **561** over link **582**, which can optionally include altering the source IP address to 40.40.40.50 as shown in FIG. 5 to provide for return traffic routing to ingress traffic node **522** instead of client node **510**. Router element **561** of VM host **560** can then receive the ingress traffic which has the updated provider network address over link **584** and determine a further destination for the ingress traffic. These further destinations can include isolated environment **570** in VM host **560**, and several network address translation elements are encountered along these pathways. These

address translations provide network isolation for software services within virtualized environments, abstract the network addressing and interfacing from physical hardware resources, and provide for enhanced migration of the software services to different physical hardware resources. Specifically, router element **561** translates the updated provider network address into a (migrated) virtual environment network address associated with an (migrated) initial isolated namespace instantiating the software service. For migrated software service **545**, the (migrated) virtual environment network address is shown as 10.0.0.4 which is associated with isolated environment **570** comprising a namespace housing further translation elements and migrated software service **545**.

(69) In one example, router element **561** then transfers the ingress traffic having the (migrated) virtual environment network address, which is received by translator element **571** of isolated environment **570**. Translator element **571** is configured to interface into the initial isolated namespace (isolated environment **570**), and translates the (migrated) virtual environment network address (10.0.0.4) to an (migrated) internal namespace network address, shown in FIG. 5 as 10.0.0.5. Translator element **571** then transfers the ingress traffic having the (migrated) internal namespace network address for receipt by another translator element **572**. In second translator element **572**, another translation occurs to translate the (migrated) internal namespace network address (10.0.0.5) to a software service network address, shown in FIG. 5 as 172.100.100.1, which is the same as pre-migration. Translator element **572** then transfers the ingress traffic having the software service network address for receipt by migrated software service **545**.

(70) Thus, a network address for software service **545** internal to an initial isolated namespace (**540**) instantiating the initial instance of software service **545** comprises a same network address for migrated software service **545** internal to a subsequent isolated namespace (**570**) instantiating the migrated instance of the software service **545**. In contrast, the provider network address comprises a different network address from the updated provider network address. This provides network migration of a network state for software service **545**, as implemented across various ingress traffic nodes, control node **520**, and various virtualized entities. Moreover, although IP addressing is discussed in FIG. 5, various TCP addressing (such as MAC addresses) can be employed along the various network links and pathways. These might also change to suit the various initial and migrated instances of the surrounding structures for software service **545**, but software service **545** itself maintains the same MAC address and IP address before, during and after migration. Any network translation for various network layers among the TCP/IP arrangement can occur among translator elements **541**, **542**, **551**, **552**, **571**, and **571**, among others.

(71) Translator elements shown in FIG. 5, such as translator element **541**, **542**, **551**, **552**, **571**, **571**, and any of the ingress traffic node translators, can be implemented as various hardware or software execution units capable of performing NAT operations, among other operations. In one example implementation, one or more of these translator elements execute as extended kernel functions executed at runtime. Examples of such kernel functions include remote procedure call (RPC) frameworks, such as the Google gRPC comprising an open-source, high-performance RPC framework developed by Google in 2015.

(72) Turning now to FIG. 7, additional migration of network connections and network properties for a software service is shown in one example implementation. Specifically, FIG. 7 illustrates migration of egress traffic routing for a software service migrating from a first instance to a second instance. Egress traffic refers to network traffic initiated and transferred by a software service for delivery to external nodes or destinations, such as other application services or nodes, client nodes, or other third-party destinations. To transfer such egress traffic, an originating software service typically includes network addressing among packets or frames to reach the destination. This network addressing, as will be discussed below, is abstracted from an actual network address of a destination, to provide for the instantiation and migration of the software service within an isolated environment, such as a network namespace within a virtualized environment or virtual machine. By

abstracting the network addressing for a software service, a destination node can continue to communicate using the same network addressing to reach a software service both before and after migration of the software service from one instance to another instance.

(73) One example of egress traffic handling is shown in system **700** of FIG. 7. System **700** includes control node **710**, egress traffic nodes **711-713**, VM host **730**, VM host **750**, and target node **770**. An explicitly labeled client node is omitted from FIG. 7, but in some examples a client node can comprise target node **770**. However, target node **770** can comprise nodes other than client nodes, such as application nodes, database nodes, third-party nodes, destination nodes, and other destinations. Target node **770** can communicate with any among egress traffic nodes **711-713** over network link **780**. Control node **710** can communicate with any among egress traffic nodes **711-713** over network link **781**. Egress traffic nodes **711-713** can communicate with various nodes, such as with VM hosts **730** and **750** over associated network links, such as network links **782-784**. VM hosts **730** and **750** can communicate over network links **785** and **786** coupled to router elements **731** and **751**. Router elements **731** and **751** can be virtualized network router elements corresponding to physical network interfaces (not shown) of host computing devices or host compute units.

(74) Target node **770** comprises various computing nodes having interface elements, network interfacing elements, data processing and storage elements, and other various components. Examples include those discussed above for client node(s) **110** above, database nodes, cloud computing nodes, authentication and authorization nodes, application service provider nodes, third-party access nodes, among others.

(75) Egress traffic nodes **711-713** comprise one or more network routing, handling, traffic processing, and transport elements or equipment. Egress traffic nodes **711-713** can include various network switching circuitry to communicatively link individual network links to other network links based on network addressing, traffic patterns, network status, or other traffic properties. In this example, egress traffic nodes **711-713** include route processor elements, such as route processor element **714**, as well as buffer elements. Route processor elements comprise network route processing (detection, evaluation, and translation) elements, including software components, data processing elements, routing table storage elements, network router elements, and other components.

(76) Control node **710** comprises control elements for instructing egress traffic nodes **711-713** and VM hosts **730** and **750** with regard to the various operational scenarios described herein. Control node **710** can include one or more microprocessors and other processing circuitry that retrieves and executes software, and any number of control software applications, from an associated storage system. Each processing element can be implemented within a single processing device but can also be distributed across multiple processing devices or sub-systems that cooperate in executing program instructions. Examples of each processing module include general purpose central processing units, application specific processors, and logic devices, as well as any other type of processing device, combinations, or variations thereof. In FIG. 7, control node **710** can take inputs from various external elements and nodes, and perform various control operations based at least on these inputs.

(77) As shown in FIG. 7, software service **745** is initially executed at a host node within isolated environment **740** as a part of a virtualized environment, e.g., VM host **730**. VM host **730** also includes router element **731** which comprises a virtualized router configured to route or otherwise handle network traffic between external network links/nodes and network elements internal to VM host **730**, such as those for isolated environment **740**. Furthermore, VM hosts **730** and **750** might further include virtualized environments within their own virtualized environments. These embedded, recursive, or nested virtualized environments can include those to instantiate software service **745**. Thus, software service **745** can comprise a further virtualized environment nested within another virtualized environment (e.g., VM hosts **730** and **750**), where the further virtualized

environment has a network interface and instantiates a client-facing software service.

(78) VM hosts **730** and **750** can be hosted by physical computing platforms, such as shown for hosts **120** and **140**, although variations are possible. Thus, these physical computing platforms might include various hardware elements including processing elements, data storage and memory elements, network interfaces, and optional user interface elements. The hardware and host software elements (e.g., host operating systems) can be isolated or virtualized from payload software elements by a virtualization system, which can include one or more virtualized environments represented by VM host **730** and **750** in FIG. 7.

(79) FIG. 7 shows various functional elements which are interconnected by network links. The actual network links are omitted from FIG. 7 for most of the functional elements, but are shown as exemplary links **780-786**. Links **780-786** can comprise any type of network link formed from one or more physical links comprising electrical, optical, or wireless signaling layers. Various network protocols and traffic types can be carried by such links, such as TCP/IP or other similar protocols associated with various network stacks and transport/internet layers. The examples herein are not intended to limit the links to any particular protocol or network link type. Also, these network links can carry various frames, packets, datagrams, or other granular data representation, which can vary based on the exact link/protocol employed. Typically, addressing used for routing the traffic is included in a header portion of a packet or frame, with payload portions configured to carry data or other packets/frames with associated addressing.

(80) Turning now to various operational scenarios for FIG. 7, operations **800** and **810** of FIG. 8 are presented. In operation **801**, an initial instance of a software service is established, such as show for software service **745** in isolated environment **740** of VM host **730**. Along with instantiating software service **745** and various storage/memory elements discussed above, corresponding network interfaces and network states can be established. These network states include various network addressing, network properties of isolated environment **740**, network routing tables, ingress/egress configurations, and other states and properties. In one example, egress traffic nodes and router elements of the VM hosts receive network address configurations from control node **710**, where the network address configurations are determined based at least on a provider network address for an initial instance of the software service. A provider or other service entity that hosts VM host **730** can provide network addressing (e.g., IP addressing) for some elements used by VM host **730**, which can correspond to physical network interfaces of a host computer system, those of router element **731**, and those associated with isolated environment **740**. Other addressing can be associated with such elements, such as media access control (MAC) addressing which typically corresponds to a physical or virtual network interface controller, and may be assigned by other entities than control node **710** or the service provider of VM host **730**.

(81) In addition to network addressing and other network properties, egress traffic nodes can be selected to service egress traffic for software service **745** and for target node **770**. The egress traffic nodes can be selected based on various factors, such as proximity (geographic or network), latency, service availability, congestion, cost, bandwidth, routing properties, or other factors. In FIG. 7, a first egress traffic node **712** is selected for software service **745** (operation **801**) and a second egress traffic node **711** is selected for target node **770** (operation **802**). In this example, egress traffic nodes for each among target node **770** and software service **745** are selected based on a lowest traffic latency between the corresponding target node or provider node (hosting software service **745**) and a particular egress traffic node. Egress traffic node **711-712** can be configured with various network routing configurations and NAT network addressing tables based on commands or instructions received from control node **710**, among other control entities.

(82) After network states have been configured/assigned for egress traffic nodes and software services within any corresponding virtualized environments, as well as egress traffic nodes selected for individual nodes, then network traffic might be issued by software service **745** for receipt by target node **770**. For example, software service **745** can originate or otherwise transfer network

traffic, such as service requests, data requests, streaming media, commands, or other traffic, over link **785** for delivery to target node **770** using an assigned network address. Moreover, this network address might be associated with a particular egress traffic node, such as egress traffic node **712**. As seen in FIG. 7, egress traffic node **712** might have a network address of 40.40.40.61 associated therewith, while egress traffic node **711** has 40.40.40.60 and egress traffic node **713** has 40.40.40.62. Software service **745** can use a network address associated with target node **770** and router element **731** can translate this network address and use 40.40.40.61 to on-ramp network traffic intended for target node **770** at egress traffic node **712**.

(83) Similar to that discussed in FIG. 5, a software service might have various network abstractions, translations (NAT or otherwise), isolations, and other elements configured to isolate a network namespace and network address of software service **745** to provide enhanced migration services. In this example, software service **745** has source IP address of 172.100.100.100 (with corresponding destination IP address of 172.100.100.1), and traffic issued by software service **745** first is delivered to router element **731** of VM host **730**, which may include various translator elements for entry/exit from isolated environment **740**, such as those noted for ingress traffic in FIG. 5. Router element **731** is configured to communicate with egress traffic node **712** over link **785** and router element **731** transfers or routes (operation **803**) this traffic for delivery to a network address associated with egress traffic node **712** (40.40.40.61), instead of an IP address of target node **770**. Thus, egress traffic node **712** receives initial egress traffic transferred by the initial instance of software service **745** for delivery to a destination node (e.g., target node **770**) through a first egress traffic node **712**.

(84) From here, egress traffic node **712** receives the egress traffic (operation **804**). Egress traffic node **712**, which includes a route processor element, can process the egress traffic issued by software service **745** to determine a route (operation **805**) to an egress traffic node associated with target node **770** and perform a NAT process to translate any associated network addressing to achieve such a route. In this example, egress traffic node **711** is associated with target node **770**, and egress traffic node **712** determines that IP address 40.40.40.60 should be used to reach egress traffic node **711**. Egress traffic node **712** updates a destination IP address of the egress traffic with the aforementioned IP address for egress traffic node **711**. Moreover, egress traffic has a return IP address updates to one associated with egress traffic node **712** (40.40.40.71), and this egress traffic is transferred for delivery (operation **806**) to egress traffic node **711** associated with target node **770** over link **783**. This egress traffic can transit various network links, networks, internetworks, internets, and other arrangements to reach egress traffic node **711** over link **782**.

(85) At egress traffic node **711**, the egress traffic can have a source network address updated using a NAT process by route processor **714** to one associated with egress traffic node **711**, such as 40.40.40.70. This egress traffic can then be transferred for delivery to target node **770** (operation **807**), which receives the egress traffic over link **780**. Target node **770** might then have return traffic in response to the egress traffic, referred to herein as egress return traffic. This egress return traffic might comprise a response to a request issued by software service **745**, return data, acknowledgement packets, heartbeat packets, new unrelated traffic, or other various network traffic. Target node **770** issues this egress return traffic using a network address associated with egress traffic node **711** (40.40.40.70) and transfers over link **780**.

(86) From here, egress traffic node **711** receives the egress return traffic. Egress traffic node **711**, which includes route processor element **714**, can process the egress return traffic issued by target node **770** to determine a route to an egress traffic node associated with software service **745** and perform a NAT process to translate any associated network addressing to achieve such a route. However, the route determination might take several forms, and the selected technique can vary in real-time or based on various criteria. The route typically transits over one or more networks, such as the Internet at-large, which is represented in FIG. 7. This route can encompass many intermediary steps, hops, routes, links, and associated equipment. In one example, at least a portion

of a network route from egress traffic node **711** to egress traffic node **712** is determined by processing a network address associated with egress traffic node **711** and/or egress traffic node **712** against a global border gate protocol (BGP) routing table for the entire Internet that is cloned or cached locally at egress traffic node **711**. This global BGP can comprise a large dataset, on the order of gigabytes (GB) of data, and can be cached into one or more storage devices of egress traffic node **711**, and can be continually updated from Internet-based nodes that house official copies of such data. A /24 CIDR address format can be employed in this example to provide such route determination. Based on processing the network addressing of the egress return traffic against this global BGP data set, a route can be determined to reach egress traffic node **712**. In another example, the egress return data can be transferred by egress traffic node **711** and employ classic Internet-based route selection, such that any links/routers along the pathway from egress traffic node **711** to egress traffic node **712** are selected en route. This second example thus employs the Internet global BGP, but as employed by various distributed and distant routing control nodes found throughout the Internet at-large. This second example is typically slower than the first example with a locally cloned global BGP data set.

(87) To continue the egress return traffic routing operations, egress traffic node **712** is associated with software service **745**, and egress traffic node **711** determines that IP address 40.40.40.71 should be used to reach egress traffic node **712**. Egress traffic node **711** updates a destination IP address of the egress traffic with the aforementioned IP address for egress traffic node **712**. Moreover, egress traffic has a return IP address updates to one associated with egress traffic node **711** (40.40.40.60), and this egress traffic is transferred for delivery (operation **811**) to egress traffic node **712** associated with a provider node hosting software service **745**. This egress traffic can transit various network links, networks, internetworks, internets, and other arrangements to reach egress traffic node **712** over link **783**.

(88) At present, the initial instance of software service **745** has not yet been migrated, and thus egress node **712** determines that no further egress routing is necessary to reach router element **731** (operation **813**, **814**, and **815**). Thus, the network traffic corresponding to the egress return traffic is delivered to router element **731**, which then routes the egress return traffic for delivery to software service **745** through various translator elements (such as those shown in FIG. 5).

(89) At some point during operation of software service **745**, a migration trigger event is produced, such as by control node **710**, and notified to various elements among system **700**, including egress traffic nodes (operation **821**). This migration trigger event then initiates various network state migration activities for a software service (as well as data/state migration and ingress traffic routing noted herein). Before or during the migration timeframe, execution of software service **745** is paused/halted at VM host **730** and software service **745** is instantiated at VM host **750** within isolated environment **760**. Various virtualized environment state/data or storage/memory data can be migrated (pre- or post-migration, noted above), and network state can be configured to support the new migrated instance at VM host **750**. However, different network addressing will be associated with the migrated instance of software service **745** at VM host **750**, as well as different network locations for VM host **750** versus VM host **730**.

(90) Control node **710** can receive this different network addressing associated with VM host **750** from a provider or other service node corresponding to VM host **750**. Control node **710** can determine an updated egress traffic node to be associated with migrated software service **745** at VM host **750**, and in some instances determines another updated egress traffic node to be associated with target node **770** (operation **822**). The different egress traffic nodes for after migration of software service **745** can be selected according to criteria noted above, such as based on traffic latency, network proximity, geographic proximity, or other factors.

(91) Control node **710** updates various nodes with this network addressing and indications of the updated/selected egress traffic nodes (operation **823**). For instance, egress traffic node **712** receives an updated network address configuration from control node **710** which indicates that migrated

software service **745** is associated with a different egress traffic node, such as egress traffic node **713**, and egress traffic node **711** receives an updated network address configuration that indicates egress traffic node **713** is to be used to reach migrated software service **745**. The updated network address configuration can be employed for determining routing of egress traffic and egress return traffic among egress traffic nodes, as well as for maintaining the same network addressing for endpoint nodes (e.g., target node **770**) to reach migrated software service **745**.

(92) Once the updated network configuration and egress traffic node selections are made/distributed, and any corresponding migration completion indication has been received/notified, then operation can resume for migrated software service **745**. However, after migration, router element **751** in VM host **750** will employ address 40.40.40.62 to reach egress traffic node **713** instead of 40.40.40.61 to reach egress traffic node **712**. In some examples, egress traffic or egress return traffic might be buffered in one or more egress traffic nodes during migration, such that after migration, a manifest of such traffic is delivered using updated addressing for the migrated instance of the software service. This can occur similar to that of the ingress traffic buffering noted above, and for a short duration of 100-200 ms. The migration completion indication denotes that the migrated instance of the software service has started at the new location/machine/provider.

(93) After migration completes, egress traffic node **713** can receive further (new) egress traffic transferred by migrated software service **745** for delivery to the same network address to reach target node **770** used before migration. Egress traffic node **713** then translates the return network address to that of egress traffic node **713**, and transfers the egress traffic for delivery to egress traffic node **711**. In addition to no noticeable interruption in service to endpoints, target node **770** and the migrated initial of software service **745** employs the same unchanged network addressing as before migration, while the egress traffic and egress return traffic are routed through one or more different egress traffic nodes after migration. Thus, the migration can be considered unrevealed and transparent to target node **770** and software service **745**. Advantageously, by selecting different egress nodes for software service **745** based on the migration, a lowest latency egress traffic node can be selected at all times, for all instances of software service and for other nodes.

(94) To further discuss egress traffic and egress return traffic after migration, the following operations can be performed, similar to that noted above for operations **800** and **810**. Network traffic might be issued by migrated software service **745** at VM host **750** for receipt by target node **770**. For example, migrated software service **745** can originate or otherwise transfer network traffic, such as service requests, data requests, streaming media, commands, or other traffic, over link **786** for delivery to target node **770** using an assigned network address. Moreover, this network address might be associated with a particular egress traffic node, such as egress traffic node **713**. As seen in FIG. 7, egress traffic node **713** might have a network address of 40.40.40.62 associated therewith, while egress traffic node **711** has 40.40.40.60. Migrated software service **745** can use a network address associated with target node **770** (same network address before migration) and router element **751** can translate this network address and use 40.40.40.62 to on-ramp network traffic intended for target node **770** at egress traffic node **713**.

(95) Migrated software service **745** might have various network abstractions, translations (NAT or otherwise), isolations, and other elements configured to isolate a network namespace and network address of migrated software service **745**. In this example, migrated software service **745** has source IP address of 172.100.100.100 (with corresponding destination IP address of 172.100.100.1), and traffic issued by migrated software service **745** first is delivered to router element **751** of VM host **750**, which may include various translator elements for entry/exit from isolated environment **760**, such as those noted for ingress traffic in FIG. 5. Router element **751** is configured to communicate with egress traffic node **713** over link **786** and router element **751** transfers or routes (operation **803**) this traffic for delivery to a network address associated with egress traffic node **713** (40.40.40.62), instead of an IP address of target node **770**. Thus, egress

traffic node **713** receives initial egress traffic transferred by the migrated instance of software service **745** for delivery to a destination node (e.g., target node **770**) through a post-migration egress traffic node **713**.

(96) From here, egress traffic node **713** receives the egress traffic (operation **804**). Egress traffic node **713**, which includes a route processor element, can process the egress traffic issued by migrated software service **745** to determine a route (operation **805**) to an egress traffic node associated with target node **770** and perform a NAT process to translate any associated network addressing to achieve such a route. In this example, egress traffic node **711** is still associated with target node **770**, and egress traffic node **713** determines that IP address 40.40.40.60 should be used to reach egress traffic node **711**. Egress traffic node **713** updates a destination IP address of the egress traffic with the aforementioned IP address for egress traffic node **711**. Moreover, egress traffic has a return IP address updated to one associated with egress traffic node **713** (40.40.40.72), and this egress traffic is transferred for delivery (operation **806**) to egress traffic node **711** associated with target node **770** over link **784**. This egress traffic can transit various network links, networks, internetworks, internets, and other arrangements to reach egress traffic node **711** over link **782**.

(97) At egress traffic node **711**, the egress traffic can have a source network address updated using a NAT process by route processor **714** to one associated with egress traffic node **711**, such as 40.40.40.70. This egress traffic can then be transferred for delivery to target node **770** (operation **807**), which receives the egress traffic over link **780**. Target node **770** might then have return traffic in response to the egress traffic, referred to herein as egress return traffic. This egress return traffic might comprise a response to a request issued by software service **745**, return data, acknowledgement packets, heartbeat packets, new unrelated traffic, or other various network traffic. Target node **770** issues this egress return traffic using a network address associated with egress traffic node **711** (40.40.40.70) and transfers over link **780**.

(98) From here, egress traffic node **711** receives the egress return traffic. Egress traffic node **711**, which includes route processor element **714**, can process the egress return traffic issued by target node **770** to determine a route to an egress traffic node associated with migrated software service **745** and perform a NAT process to translate any associated network addressing to achieve such a route. However, the route determination might take several forms, and the selected technique can vary in real-time or based on various criteria. The route typically transits over one or more networks, such as the Internet at-large, which is represented in FIG. 7. This route can encompass many intermediary steps, hops, routes, links, and associated equipment. In one example, at least a portion of a network route from egress traffic node **711** to egress traffic node **713** is determined by processing a network address associated with egress traffic node **711** and/or egress traffic node **713** against a global border gate protocol (BGP) routing table for the entire Internet that is cloned or cached locally at egress traffic node **711**. This global BGP can comprise a large dataset, on the order of gigabytes (GB) of data, and can be cached into one or more storage devices of egress traffic node **711**, and can be continually updated from Internet-based nodes that house official copies of such data. A /24 CIDR address format can be employed in this example to provide such route determination. Based on processing the network addressing of the egress return traffic against this global BGP data set, a route can be determined to reach egress traffic node **713**. In another example, the egress return data can be transferred by egress traffic node **711** and employ classic Internet-based route selection, such that any links/routers along the pathway from egress traffic node **711** to egress traffic node **713** are selected en route. This second example thus employs the Internet global BGP, but as employed by various distributed and distant routing control nodes found throughout the Internet at-large. This second example is typically slower than the first example with a locally cloned global BGP data set.

(99) To continue the egress return traffic routing operations, egress traffic node **713** is associated with migrated software service **745**, and egress traffic node **711** determines that IP address

40.40.40.72 should be used to reach egress traffic node **713**. Egress traffic node **711** updates a destination IP address of the egress traffic with the aforementioned IP address for egress traffic node **713**. Moreover, egress traffic has a return IP address updates to one associated with egress traffic node **711** (40.40.40.60), and this egress traffic is transferred for delivery (operation **811**) to egress traffic node **713** associated with a provider node hosting migrated software service **745**. This egress traffic can transit various network links, networks, internetworks, internets, and other arrangements to reach egress traffic node **713** over link **784**.

(100) FIG. **9** illustrates node system **900** and associated software **905** in an implementation. FIG. **9** illustrates node system **900** that is representative of any system or collection of systems in which the various operational architectures, scenarios, and processes disclosed herein may be implemented. For example, node system **900** can be used to implement client nodes, provider nodes (such as VM hosts, servers, hardware for virtualized environments), control nodes, ingress traffic nodes, egress traffic nodes, and various other endpoint, control, and routing nodes discussed herein. Specifically, client nodes **110**, network traffic nodes **190**, and host nodes **120** and **140** of FIG. **1** might be implemented by elements of node system **900**. Client nodes **310**, network traffic nodes **390**, host nodes **320** and **340**, and data storage system **360** of FIG. **3** might be implemented by elements of node system **900**. Client node **510**, control node **520**, ingress traffic nodes **521-523**, and host machines for VM hosts **530** and **560** of FIG. **5** might be implemented by elements of node system **900**. Control node **710**, egress traffic nodes **711-713**, host machines for VM hosts **730** and **750**, and target node **770** of FIG. **7** might be implemented by elements of node system **900**.

Variations on these implementations are possible, with some examples only employing a selected portion of the elements discussed for node system **900**.

(101) Node system **900** may be implemented as a single apparatus, system, or device or may be implemented in a distributed manner as multiple apparatuses, systems, or devices. Node system **900** includes, but is not limited to, processing system **902**, storage system **903**, software **905**, communication interface system **907**, user interface system **908**, and router system **909**. Processing system **902** is operatively coupled with storage system **903**, communication interface system **907**, user interface system **908**, and router system **909**.

(102) Processing system **902** loads and executes software **905** from storage system **903**. Software **905** includes applications **920**, which are representative of the processes, services, and platforms discussed with respect to the included Figures. Various portion of software **905** might be included or excluded based on what type of node or system is implemented by node system **900**. Thus, all software elements might not be present in every implementation of node system **900**. When executed by processing system **902** software **905** directs processing system **902** to operate as described herein for at least the various processes, operational scenarios, and sequences discussed in the foregoing implementations. Node system **900** may optionally include additional devices, features, or functionality not discussed for purposes of brevity.

(103) Referring still to FIG. **9**, processing system **902** may comprise a microprocessor and processing circuitry that retrieves and executes software **905** from storage system **903**. Processing system **902** may be implemented within a single processing device but may also be distributed across multiple processing devices or sub-systems that cooperate in executing program instructions. Examples of processing system **902** include general purpose central processing units, application specific processors, and logic devices, as well as any other type of processing device, combinations, or variations thereof.

(104) Storage system **903** may comprise any computer readable storage media readable by processing system **902** and capable of storing software **905**. Storage system **903** may include volatile and nonvolatile, removable and non-removable media implemented in any method or technology for storage of information, such as computer readable instructions, data structures, program modules, or other data. Examples of storage media include random access memory, read only memory, magnetic disks, optical disks, flash memory, virtual memory and non-virtual

memory, magnetic memory, magnetic storage devices, or any other suitable storage media. In no case is the computer readable storage media a propagated signal. In addition to computer readable storage media, in some implementations storage system **903** may also include computer readable communication media over which at least some of software **905** may be communicated internally or externally. Storage system **903** may be implemented as a single storage device but may also be implemented across multiple storage devices or sub-systems co-located or distributed relative to each other. Storage system **903** may comprise additional elements, such as a controller, capable of communicating with processing system **902** or possibly other systems.

(105) Software **905** may be implemented in program instructions and among other functions may, when executed by processing system **902**, direct processing system **902** to operate as described with respect to the various operational scenarios, sequences, and processes illustrated herein. For example, software **905** may include program instructions comprising applications **920**, operating system **921**, and data **922**. In particular, the program instructions may include various components or modules that cooperate or otherwise interact to carry out the various processes and operational scenarios described herein. The various components or modules may be implemented in compiled or interpreted instructions, or in some other variation or combination of instructions. The various components or modules may be executed in a synchronous or asynchronous manner, serially or in parallel, in a single threaded environment or multi-threaded, or in accordance with any other suitable execution paradigm, variation, or combination thereof. Software **905** may include additional processes, programs, or components, such as operating system software or other application software, in addition to or that include applications **920**. Software **905** may also comprise firmware or some other form of machine-readable processing instructions executable by processing system **902**.

(106) Software **905**, when loaded into processing system **902** and executed, may transform a suitable apparatus, system, or device (of which node system **900** is representative) overall from a general-purpose computing system into a special-purpose computing system customized to implement client nodes, provider nodes, control nodes, ingress traffic nodes, egress traffic nodes, and various other endpoint, control, and routing nodes discussed herein. Indeed, encoding software **905** on storage system **903** may transform the physical structure of storage system **903**. The specific transformation of the physical structure may depend on various factors in different implementations of this description. Examples of such factors may include, but are not limited to, the technology used to implement the storage media of storage system **903** and whether the computer-storage media are characterized as primary or secondary storage, as well as other factors. For example, if the computer-readable storage media are implemented as semiconductor-based memory, software **905** may transform the physical state of the semiconductor memory when the program instructions are encoded therein, such as by transforming the state of transistors, capacitors, or other discrete circuit elements constituting the semiconductor memory. A similar transformation may occur with respect to magnetic or optical media. Other transformations of physical media are possible without departing from the scope of the present description, with the foregoing examples provided only to facilitate the present discussion.

(107) Applications **920** can include routing control system **930**, migration control system **935**, and virtualization system **940**. Routing control system **930** includes address translator **931**, ingress/egress node selector **932**, and route selector **933**. Migration control system **935** includes migration initiator **936**, migration data/state handler **937**, and migration network handler **938**. Virtualization system **940** includes virtual machine managers **941**, network interface **942**, isolated environment manager **943**, and application/service instances **944**.

(108) Turning first to routing control system **930**, address translator **931** performs various network address translation and packet/frame updating for altering network addressing associated with received and transferred network traffic. This can include NAT functions using routing data **945** or translation tables **947** in data **922**. Ingress/egress node selector **932** can include operations to select

specific ingress or egress traffic nodes according to various criteria noted herein to provide ingress or egress of network traffic for endpoint nodes including client nodes, software services, and other elements. Ingress/egress node selector **932** can select nodes responsive to migration triggers or to initial instantiations of software services, or according to various changes in network conditions, among other factors. Route selector **933** can provide route selection across one or more networks to route egress or ingress traffic among pairs of egress nodes or pairs of ingress nodes, among other nodes. Route selector **933** might use the routing of external routers, such as the Internet, or may determine routes based on locally cached versions of global routing tables or global BGP tables stored in data **922**.

(109) Turning next to migration control system **935**, migration initiator **936** includes detection routines that determine when a software service is to be migrated from an initial instance to a migrated instance. Triggers for this migration can include various criteria included in migration parameters **948**, such as network conditions, network congestion, provider costs, bandwidth costs, provider availability changes, altered locality preferences of clients, time of day, day/month, provider policy changes or terms, or other factors discussed herein. Migration data/state handler **937** determines which data is to be migrated for software services, such as done for the pre/post migration techniques discussed herein, can establish various mappings or block devices that implement memory/storage devices of virtualized software services, and can perform copying or migration of data/state during migration operations. Migration network handler **938** performs migration of network state and network addressing for migration of software services, including determination of new network addressing provided by new providers/hosts, updating of ingress/egress traffic nodes with network addressing before and at migration, and indications to router elements of VM hosts of network addressing changes for instantiation and migration of software services.

(110) Turning next to virtualization system **940** virtual machine managers **941** instantiate one or more virtualized environments that host software services, such as VM hosts, to perform the functions of the software services on the associated host hardware or machines. Virtual machine managers **941** can establish router elements (through network interface **942**) which communicate with isolated environments or network namespaces for software services, and establish instances of software services within such isolated environments or network namespaces. Isolated environment manager **943** can work with virtual machine managers **941** to establish isolated environments or network namespaces, as well as establish translator elements to translate network traffic addressing among various network address spaces such that software services can be instantiated with the same network addressing for each instance. Application/service instances **944** include actual instantiations of software services, which can include virtualized instantiations, and instantiations of initial and migrated software services.

(111) Communication interface system **907** may include communication connections and devices that allow for communication with other computing systems or electrical components (not shown) over communication links or communication networks (not shown). Examples of connections and devices that together allow for inter-system communication may include network interface controllers, network interfacing circuitry, transceivers, antennas, power amplifiers, RF circuitry, and other communication circuitry. The connections and devices may communicate over communication media to exchange communications with other computing systems or networks of systems, such as metal, glass, air, or any other suitable communication media. Physical or logical elements of communication interface system **907** can provide network information, router information, and other information. Communication interface system **907** may include portions of router system **909** or communication interfaces for communicating with router equipment.

(112) Communication between communication node system **900** and other elements or systems (not shown), may occur over communication links or communication networks and in accordance with various communication protocols, combinations of protocols, or variations thereof provided

by at least communication interface system **907**. For example, communication node system **900** when implementing a control device, might communicate with sensor elements over corresponding digital communication links comprising Ethernet interfaces, serial interfaces, serial peripheral interface (SPI) links, inter-integrated circuit (I2C) interfaces, universal serial bus (USB) interfaces, UART interfaces, or wireless interfaces. When network links are employed, example networks include intranets, internets, the Internet, local area networks, wide area networks, wireless networks, wired networks, virtual networks, software defined networks, data center buses, computing backplanes, or any other type of network, combination of network, or variation thereof. The aforementioned communication networks and protocols are well known and need not be discussed at length here. However, some network communication protocols that may be used include, but are not limited to, the Ethernet, Internet protocol (IP, IPv4, IPv6, etc. . . .), the transmission control protocol (TCP), and the user datagram protocol (UDP), as well as any other suitable communication protocol, variation, or combination thereof.

(113) User interface system **908** is optional in some implementations, and may include a software or virtual interface such as a terminal interface, command line interface, or application programming interface (API). User interface system **908** may also include physical user interfaces, such as keyboard, a mouse, a voice input device, or a touchscreen input device for receiving input from a user. User interface system **908** may include telemetry interfaces, router status interfaces, user command controls, router operation mode command controls, and user interface indications, visualizations, and representations, among others. Output devices such as displays, web interfaces, terminal interfaces, and other types of output devices may also be included in user interface system **908**. User interface system **908** can provide output and receive input over a network interface, such as communication interface system **907**. In network examples, user interface system **908** might packetize data for receipt by a display system or computing system coupled over one or more network interfaces. User interface system **908** may comprise API elements for interfacing with users, other data systems, other user devices, web interfaces, and the like. User interface system **908** may also include associated user interface software executable by processing system **902** in support of the various user input and output devices discussed above. Separately or in conjunction with each other and other hardware and software elements, the user interface software and user interface devices may support a console user interface, graphical user interface, a natural user interface, or any other type of user interface.

(114) Router system **909** comprises various hardware and software elements for interfacing with router equipment, or alternatively, router equipment configured to route network traffic. Router system **909** includes one or more network interfaces, including network interface controller circuitry capable of receiving and transferring network traffic employed for corresponding network links and various layered protocols. Router system **909** can comprise network traffic translation elements configured to perform NAT functions, which can include portions of software **905**. Moreover, router system **909** can comprise one or more software components which execute via processing system **902** to perform network traffic routing, translation, buffering, ingress/egress handling, and other operations. Router system **909** can comprise elements of communication interface system **907**, such as those for physically interfacing with network links (e.g., Ethernet or various IP-compatible links). Router system **909** can also perform various modification of network traffic, including network frames or network packets, to modify addressing of header portions, perform deep packet inspection for routing optimization, and other functions. These various operations and functions can be performed by software **905** executed by processing system **902**, or may include dedicated processing hardware and interface circuitry.

(115) The functional block diagrams, operational scenarios and sequences, and flow diagrams provided in the Figures are representative of exemplary systems, environments, and methodologies for performing novel aspects of the disclosure. While, for purposes of simplicity of explanation, methods included herein may be in the form of a functional diagram, operational scenario or

sequence, or flow diagram, and may be described as a series of acts, it is to be understood and appreciated that the methods are not limited by the order of acts, as some acts may, in accordance therewith, occur in a different order and/or concurrently with other acts from that shown and described herein. For example, those skilled in the art will understand and appreciate that a method could alternatively be represented as a series of interrelated states or events, such as in a state diagram. Moreover, not all acts illustrated in a methodology may be required for a novel implementation.

(116) Furthermore, it should be understood that the disclosures and enhancements herein are applicable across a range of suitable systems, elements, and operations. For example, network links and network protocols can vary based on application, and may evolve according to changes and adoption of different standards. Thus, the descriptions and figures included herein depict specific implementations to teach those skilled in the art how to make and use the best options. For the purpose of teaching inventive principles, some conventional aspects have been simplified or omitted. Those skilled in the art will appreciate variations from these implementations that fall within the scope of this disclosure. Those skilled in the art will also appreciate that the features described above can be combined in various ways to form multiple implementations.

Claims

1. A method, comprising: in a router element for an initial instance of a software service, receiving egress traffic initiated by the software service for delivery to a destination node, the egress traffic indicating a destination address different from an actual destination address of the destination node, and adjusting the destination address of the egress traffic to correspond to a first egress traffic node assigned to service the initial instance of the software service; in the first egress traffic node, altering the destination address of the egress traffic to indicate a second egress traffic node associated with the destination node, and transferring the egress traffic for delivery to the second egress traffic node that transfers the egress traffic for delivery to the destination node using the actual destination address; in the second egress traffic node, receiving return traffic transferred by the destination node for delivery to the initial instance of the software service, altering a return traffic destination address to indicate the first egress traffic node, and transferring the return traffic for delivery to the first egress traffic node that transfers the return traffic for delivery to the router element for the initial instance of the software service; responsive to migration of the software service to a subsequent instance, assigning a third egress traffic node to the migrated instance of the software service, wherein the third egress traffic node is configured to receive subsequent egress traffic for delivery to the destination node transferred by a router element for the subsequent instance of the software service.
2. The method of claim 1, comprising: in the second egress traffic node, receiving the return traffic transferred by the destination node for delivery to the subsequent instance of the software service, altering the return traffic destination address to indicate the third egress traffic node, and transferring the return traffic for delivery to the third egress traffic node that transfers the return traffic for delivery to the router element for the subsequent instance of the software service.
3. The method of claim 1, wherein selection of the first egress traffic node and the third egress traffic node is based at least on traffic latency with respect to the software service at the initial instance and the subsequent instance.
4. The method of claim 1, comprising: selecting at least a portion of a network route for the return traffic from the second egress traffic node to the first egress traffic node by at least processing a network address associated with the first egress traffic node against a global border gate protocol (BGP) routing table stored at the second egress traffic node.
5. The method of claim 1, wherein the software service is assigned a service network address within an initial isolated namespace which remains the same after migration to a subsequent initial

isolated namespace.

6. The method of claim 1, comprising: in a control node, indicating network translations to the first egress traffic node, the second egress traffic node, and the third egress traffic node for handling the egress traffic.

7. The method of claim 1, comprising: in a control node, responsive to migration of the software service to the subsequent instance, indicating to the first egress traffic node that the first egress traffic node no longer handles traffic for the subsequent instance of the software service, and indicating to the second egress traffic node that the third egress traffic node replaces the second egress traffic node for routing traffic associated with the software service.

8. The method of claim 1, comprising: responsive to a migration trigger event for migration of the software service to the subsequent instance, in the second egress traffic node, buffering the return traffic addressed for transfer to the first egress traffic node for delivery to the initial instance of the software service; responsive to receipt of a migration completion indication, in the second egress traffic node, altering a destination address of the return traffic to indicate the third egress traffic node, and transferring the buffered return traffic for delivery to the third egress traffic node that transfers the buffered return traffic for delivery to the router element for the subsequent instance of the software service.

9. The method of claim 1, wherein preceding the migration, the egress traffic is routed by at least: receiving, in a translator element of an initial isolated namespace instantiating the software service, the egress traffic having a first destination address from the software service, and translating the first destination address to a second destination address associated with the router element of the initial instance; receiving, in the router element of the initial instance, the egress traffic having the second destination address, and translating the second destination address into a third destination network address associated with the first egress traffic node.

10. The method of claim 9, wherein after the migration, the subsequent egress traffic is routed by at least: receiving, in a translator element of a subsequent isolated namespace instantiating the software service, the subsequent egress traffic having a fourth destination address from the software service, and translating the fourth destination address to a fifth destination address associated with the router element of the subsequent instance; receiving, in the router element of the subsequent instance, the subsequent egress traffic having the fifth destination address, and translating the fifth destination address into a sixth destination network address associated with the third egress traffic node.

11. A system, comprising: a router element for an initial instance of a software service receives egress traffic initiated by the software service for delivery to a destination node, the egress traffic indicating a destination address different from an actual destination address of the destination node, and adjust the destination address of the egress traffic to correspond to a first egress traffic node assigned to service the initial instance of the software service; the first egress traffic node alters the destination address of the egress traffic to indicate a second egress traffic node associated with the destination node, and transfers the egress traffic for delivery to the second egress traffic node that transfers the egress traffic for delivery to the destination node using the actual destination address; the second egress traffic node receives return traffic transferred by the destination node for delivery to the initial instance of the software service, alters a return traffic destination address to indicate the first egress traffic node, and transfers the return traffic for delivery to the first egress traffic node that transfers the return traffic for delivery to the router element for the initial instance of the software service; and responsive to migration of the software service to a subsequent instance, a control node assigns a third egress traffic node to the migrated instance of the software service, wherein the third egress traffic node receives subsequent egress traffic for delivery to the destination node transferred by a router element for the subsequent instance of the software service.

12. The system of claim 11, comprising: the second egress traffic node receives the return traffic transferred by the destination node for delivery to the subsequent instance of the software service,

alters the return traffic destination address to indicate the third egress traffic node, and transfers the return traffic for delivery to the third egress traffic node that transfers the return traffic for delivery to the router element for the subsequent instance of the software service.

13. The system of claim 11, wherein selection of the first egress traffic node and the third egress traffic node is based at least on traffic latency with respect to the software service at the initial instance and the subsequent instance.

14. The system of claim 11, comprising: the second egress traffic node selects at least a portion of a network route for the return traffic from the second egress traffic node to the first egress traffic node by at least processing a network address associated with the first egress traffic node against a global border gate protocol (BGP) routing table stored at the second egress traffic node.

15. The system of claim 11, wherein the software service is assigned a service network address within an initial isolated namespace which remains the same after migration to a subsequent initial isolated namespace.

16. The system of claim 11, comprising: the second egress traffic node, responsive to a migration trigger event for migration of the software service to the subsequent instance, buffers the return traffic addressed for transfer to the first egress traffic node for delivery to the initial instance of the software service; the second egress traffic node, responsive to receipt of a migration completion indication, alters a destination address of the return traffic to indicate the third egress traffic node, and transfers the buffered return traffic for delivery to the third egress traffic node that transfers the buffered return traffic for delivery to the router element for the subsequent instance of the software service.

17. The system of claim 11, wherein preceding the migration: a translator element of an initial isolated namespace instantiating the software service receives the egress traffic having a first destination address from the software service, and translates the first destination address to a second destination address associated with the router element of the initial instance; the router element of the initial instance receives the egress traffic having the second destination address, and translates the second destination address into a third destination network address associated with the first egress traffic node.

18. The system of claim 17, wherein after the migration: a translator element of a subsequent isolated namespace instantiating the software service receives the subsequent egress traffic having a fourth destination address from the software service, and translates the fourth destination address to a fifth destination address associated with the router element of the subsequent instance; the router element of the subsequent instance receives the subsequent egress traffic having the fifth destination address, and translates the fifth destination address into a sixth destination network address associated with the third egress traffic node.
