

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12418485
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Valancius; Vytautas et al.

Traffic shaping offload on a network interface controller

Abstract

Systems, methods, and apparatuses, including network interface controllers and computer-readable media, for traffic shaping offload. A network computing device can receive data packets for transmission and implement a traffic policy that includes transmitting at least some data packets without delay to their intended destination. Confirmation tokens for non-delayed packets can be queued in a time-indexed data structure and dequeued according to a traffic shaping policy. Confirmation tokens can be generated and stored independent of the time at which corresponding packets for the tokens are transmitted. Dequeued confirmation tokens can cause the network computing device to receive additional packets for transmission. The device can flag at least some packets for transmission without delay according to aspects of the disclosure, while un-flagged packets can be shaped according to a different traffic shaping policy.

Inventors: Valancius; Vytautas (Palo Alto, CA), Lane; Rich (Mountain View, CA), Forrest; Bailey (Sunnyvale, CA), Li; Chonggang (San Jose, CA)

Applicant: Google LLC (Mountain View, CA)

Family ID: 1000008819936

Assignee: Google LLC (Mountain View, CA)

Appl. No.: 18/233919

Filed: August 15, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230396551 A1	Dec. 07, 2023

Related U.S. Application Data

continuation parent-doc US 17114694 20201208 US 11784933 child-doc US 18233919

Publication Classification

Int. Cl.: H04L47/22 (20220101); H04L47/20 (20220101); H04L47/2441 (20220101);
H04L47/28 (20220101); H04L47/30 (20220101)

U.S. Cl.:

CPC H04L47/22 (20130101); H04L47/20 (20130101); H04L47/2441 (20130101);
H04L47/28 (20130101); H04L47/30 (20130101);

Field of Classification Search

CPC: H04L (47/22); H04L (47/20); H04L (47/2441); H04L (47/28); H04L (47/30)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6215769	12/2000	Ghani	709/234	H04L 1/1854
6925055	12/2004	Erimli et al.	N/A	N/A
7106693	12/2005	Turner et al.	N/A	N/A
7243177	12/2006	Davis et al.	N/A	N/A
7724740	12/2009	Wang et al.	N/A	N/A
7818460	12/2009	Kim et al.	N/A	N/A
8156265	12/2011	Buick et al.	N/A	N/A
9014198	12/2014	Breslin et al.	N/A	N/A
9148485	12/2014	Mannava et al.	N/A	N/A
2003/0152096	12/2002	Chapman	N/A	N/A
2007/0019550	12/2006	Enomoto	N/A	N/A
2008/0311885	12/2007	Dawson et al.	N/A	N/A
2009/0204723	12/2008	Tonsing et al.	N/A	N/A
2009/0285095	12/2008	Baden et al.	N/A	N/A
2013/0077486	12/2012	Keith	N/A	N/A
2013/0286835	12/2012	Plamondon et al.	N/A	N/A
2018/0212886	12/2017	Contavalli	N/A	H04L 47/225

OTHER PUBLICATIONS

Extended European Search Report for European Patent Application No. 21171383.9 dated Aug. 13, 2021. 10 pages. cited by applicant

Office Action for European Patent Application No. 21171383.9 dated Jul. 25, 2023. 8 pages. cited by applicant

Primary Examiner: Park; Jung H

Attorney, Agent or Firm: Lerner David LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application is a continuation of U.S. patent application Ser. No. 17/114,694, filed Dec. 8, 2020, which is now U.S. Pat. No. 11,784,933, the disclosure of which is incorporated herein by reference.

BACKGROUND

(1) In a network communicating data across one or more nodes, a variety of different techniques can be applied for managing how different packets of data are sent across the network to the various nodes. Managing the volume and transmission speed of communicated data remains a problem for a network of nodes, particularly when the various nodes are computing devices that have physical limits for transmitting data and can be quickly overwhelmed by a network if the limits are not respected.

(2) Techniques for managing network traffic include techniques for prioritizing different packets according to different criteria, by delaying or expediting transmission to and from a node in accordance with a schedule. One such class of techniques for managing network traffic can include traffic shaping. Traffic shaping can refer to modifying the rate at which different packets—or portions of packets—of data are transmitted from a node. Under a traffic-shaping policy, packets can be managed according to different queues and algorithms for determining in what order packets leave a node.

(3) Traffic shaping can also include rate-limiting, such as enforcing the rate at which packets leave or enter a node; and pacing, such as injecting inter-packet gaps to smooth traffic over time. Traffic shaping techniques can be used to manage network congestion at one or more network nodes, and can be implemented with respect to the physical limits—for example RAM or available processing power—of the different nodes of the network.

BRIEF SUMMARY

(4) The present disclosure provides for a system, method, and apparatus, including non-transitory computer-readable media, for shaping traffic out of a network computing device by transmitting at least some packets without delay through a network interface controller (NIC), while queueing confirmation tokens for each packet in a time-indexed data structure. Instead of packets being stored in memory coupled to a network interface controller, confirmation tokens are enqueued and at least some packets are sent immediately to a destination node without delay. The confirmation token for a packet is generated and enqueued independent of whether or not the corresponding data packet was actually transmitted. Confirmation tokens are dequeued from the time-indexed data structure according to a traffic shaping policy, and passed to the network protocol stack of a host network device for fetching additional data packets.

(5) Memory consumption can be reduced by storing confirmation tokens on the NIC instead of entire data packets. A NIC implemented according to aspects of this disclosure can send data packets independently of confirmation tokens to a network protocol stack, for example by sending a packet before a confirmation token is sent to the network protocol stack, or sending a confirmation token before sending a packet to a destination node. A NIC implemented in accordance with this disclosure can efficiently shape traffic without packet loss or head-of-line blocking. Head-of-line blocking can occur when a first packet is blocked, causing a second packet that is received after the first packet to be delayed. Aspects of this disclosure can be implemented in addition to other traffic shaping techniques, for example by not delaying certain flagged packets and enqueueing confirmation tokens only for the flagged packets.

(6) In general, one innovative aspect of the subject matter described in this specification provides for a network computing device including one or more processors. The one or more processors can be configured to receive data packets for transmission to one or more network nodes of a network according to a traffic shaping policy. The data packets can include a first data packet. The one or more processors can be further configured to: generate and store a first confirmation token for the first data packet in a time-indexed data structure; transmit the first data packet to the one or more

network nodes; and dequeue the first confirmation token for the first data packet, wherein dequeuing the first confirmation token causes the network computing device to receive an additional data packet for transmission.

(7) The one or more processors can be further configured to generate and store the first confirmation token before transmitting the data packet to the one or more network nodes.

(8) The received data packets can include a second data packet. The one or more processors can be further configured to transmit the second data packet to the one or more network nodes; and after the second data packet is transmitted, generate and store a second confirmation token for the second data packet in the time-indexed data structure.

(9) The one or more processors can be further configured to: flag one or more data packets of the received data packets for transmission without delay to the one or more network nodes; transmit the one or more flagged data packets to the one or more network nodes; generate and store respective confirmation tokens for the one or more flagged data packets in the time-indexed data structure; and delay the transmission of an un-flagged data packet of the received data packets by storing the un-flagged data packet at a first time and transmitting the stored un-flagged data packet after the first time.

(10) The time-indexed data structure can be a first time-indexed data structure and the traffic shaping policy can be a first traffic shaping policy. The one or more processors can be further configured to: dequeue the first confirmation token for the first data packet in accordance with the first traffic shaping policy, store the un-flagged data packet in memory coupled to the network computing device, generate and store an identifier for the un-flagged data packet in a second time-indexed data structure; and dequeue the identifier in the second time-indexed data structure in accordance with a second traffic shaping policy, wherein dequeuing the identifier causes the network computing device to transmit the un-flagged data packet in response to dequeuing the identifier in the second time-indexed data structure.

(11) The one or more flagged data packets can be one or more flagged first data packets. The one or more processors can be further configured to: identify a second data packet whose transmission has been delayed, and flag one or more third data packets for transmission without delay, wherein the one or more third data packets are received after the second data packet.

(12) The one or more processors can further configured to: maintain a current time, and generate a timestamp for the first confirmation token; and wherein the traffic shaping policy specifies dequeuing the first confirmation token when the timestamp meets or exceeds the current time.

(13) The received data packets can be received from a plurality of applications running on the network computing device. The one or more processors can be further configured to dequeue the first confirmation token according to the traffic shaping policy, wherein the traffic shaping policy prioritizes transmission of data packets from a first application of the plurality of applications over data packets from a second application of the plurality of applications.

(14) The plurality of applications running on the network computing device can include at least one application executing on a virtual machine running on the network computing device.

(15) The one or more processors can be coupled to a network interface controller, and the one or more processors can be further configured to transmit the first data packet to the one or more network nodes using the network interface controller.

(16) The network interface controller can be a programmable network interface controller.

(17) Other implementations of the foregoing aspect can include a system of one or more computing devices, a computer-implemented method, an apparatus, and computer programs recorded on one or more computer-readable storage media.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram of an example network that includes a network computing device implementing a traffic shaping engine.
- (2) FIG. 2 is a block diagram of the network computing device and the network nodes of the network.
- (3) FIG. 3 is a relational block diagram illustrating the traffic shaping engine shaping outbound data packets.
- (4) FIG. 4 is a block diagram of the time-indexed data structure for the traffic shaping engine.
- (5) FIG. 5 is a block diagram of the traffic shaping engine implementing separate time-indexed data structures for delayed and un-delayed outbound data packets.
- (6) FIG. 6 is a flowchart of an example process for traffic shaping on a NIC.
- (7) Like numbers represent like elements in the figures.

DETAILED DESCRIPTION

- (8) FIG. 1 is a block diagram of an example network **125** that includes a network computing device **110** implementing a traffic shaping engine **100**. The network computing device **110** can be coupled to a computing platform **105** through a network **125**. The network computing device **110** can be any one of a variety of devices, such as a laptop, desktop computer, a mobile device like a phone or a tablet, or a device intended to be worn, such as a watch, earbuds, glasses, or a helmet. For ease of description, reference will be made to the network computing device **110** as interacting with the computing platform **105**.
- (9) It should be appreciated that a typical computing platform can include a large number of connected server computing devices, many storage devices, and/or many datacenters located at one or more physical sites or locations. Although described in the context of the computing platform **105**, the network computing device **110** and network nodes **112A-N** can be part of any network, such as a distributed network. The computing platform **105** can include one or more datacenters **115A-N**, which can be located in the same or different physical locations. Each datacenter can implement one or more network nodes. For example, the datacenter A **115A** can include network nodes **112A-N**.
- (10) The network **125** can be arranged according to various organizational network structures or topologies. For example, the network computing device **110** and the network nodes **112A-N** can be part of a peer-to-peer network. As other examples, the network **125** can connect the network computing device **110** and the network nodes **112A-N** according to a point-to-point network topology, or a hub-and-spoke network topology with one or more central—or “hub”—devices.
- (11) A network node can include one or more discrete units of physical or virtual computing resources assigned to perform a service or application. The discrete units can be specified on the computing platform as physical and/or virtual computing resources. As a physical computing resource, a node can include one or more processors and/or one or more storage devices across the one or more datacenters **115A-N** of the computing platform **105**. As a virtual computing resource, a node can include one or more virtual machines (VM) **155A-N**, each virtual machine operating using physical processors and storage devices of the computing platform **105**.
- (12) The one or more virtual machines can provide a level of abstraction between physical computing resources and a platform user of the platform, and multiple virtual machines for different platform users can operate on shared physical computing resources. From the perspective of the platform user, however, the node can appear as a dedicated physical set of hardware, such as if the platform user were executing applications or other programs directly on the hardware.
- (13) An application can be a collection of services and supporting software, such as a mail server managing incoming and outgoing email. A service can be a specialized program that performs one or more operations as part of a running application. For example, a service for a mail server can be a service for reading incoming mail according to a particular protocol, such as POP3 or IMAP. In

addition to the network nodes **112A-N**, the network computing device **110** can also run one or more applications, and be implemented on physical and/or virtual computing resources.

(14) The computing platform **105** can provide an interface between the network computing device **110** and the computing platform **105**. The interface can be implemented in a number of ways. For example, the computing platform **105** can implement a user interface that is displayed on a display connected to the network computing device **110**.

(15) The network computing device **110** and the network nodes **112A-N** can be configured to communicate over the network **125**. The network **125** itself can include various configurations and protocols including the Internet, World Wide Web, intranets, virtual private networks, wide area networks, local networks, and private networks using communication protocols proprietary to one or more companies. The network **125** can support a variety of short- and long-range connections along a variety of different bandwidths, such as 2.402 GHz to 2.480 GHz commonly associated with the Bluetooth® standard, 2.4 GHz and 5 GHz commonly associated with the Wi-Fi® communication protocol, or with a variety of communication standards, such as the LTE® standard for wireless broadband communication. The network **125** can also support wired connections between the network computing device **110**, and the network nodes **112A-N**, including over various types of Ethernet connection.

(16) The network computing device **110** can implement the traffic shaping engine **100**. As described in more detail herein, the traffic shaping engine **100** is configured to shape traffic transmitted to and from the network computing device **110**. The traffic shaping engine **100** can be implemented as part of a network interface controller (NIC) **220**. The NIC **220** is a component of the network computing device **110** for communicating data to and from the network computing device **110** and network nodes of the network **125**. The NIC **220** can be implemented as a combination of hardware and software. The NIC **220** can include hardware such as integrated circuits for physically sending data through the network **125**. For example, the NIC **220** can be implemented as a card that slots into a bus of a motherboard, or the NIC **220** can be integrated as part of a motherboard or other circuit board of the network computing device **110**.

(17) The NIC **220** can implement any of a variety of different features, including features for shaping packets of data to and from the network computing device **110**. Traffic shaping features can be advantageously implemented on NICs of nodes as opposed to being implemented elsewhere in the network **125**, such as on middle devices linking nodes of the network **125**. The burden of computational processing required of traffic shaping can be distributed across the various devices in communication over the network, as opposed to relying on a predetermined number of devices connected elsewhere on the network **125**.

(18) The traffic shaping engine **100** can be implemented in any of a variety of different ways in accordance with aspects of this disclosure. For example, the traffic shaping engine **100** can be one or more circuits physically implemented on the NIC **220**. Alternatively or in addition, the traffic shaping engine **100** can be programmatically implemented as part of the NIC **220**. In some implementations in which the NIC **220** is programmable, the traffic shaping engine **100** can be installed as a program on the NIC **220** that is executed when the NIC **220** is in operation.

(19) FIG. 2 is a block diagram of the network computing device **110** and the network nodes **112A-N** of the network **125**. The storage device(s) **230** can be a combination of volatile and non-volatile memory, and can be at the same or different physical locations as the network computing device **110** and the network nodes **112A-N**. For example, the storage device(s) **230** can include any type of non-transitory computer readable medium capable of storing information, such as a hard-drive, solid state drive, tape drive, optical storage, memory card, ROM, RAM, DVD, CD-ROM, write-capable, and read-only memories

(20) The network computing device **110** can include one or more processors **212** and memory **214**. The memory **214** can store information accessible by the processor(s) **212**, including instructions **216** that can be executed by the processors **212**. The memory **214** can also include data **218** that

can be retrieved, manipulated or stored by the processor(s) **212**. The memory **214** can be a type of non-transitory computer readable medium capable of storing information accessible by the processors **212**, such as volatile and non-volatile memory. The processor(s) **212** can include one or more central processing units (CPUs), graphic processing units (GPUs), field-programmable gate arrays (FPGAs), and application-specific integrated circuits (ASICs), such as tensor processing units (TPUs).

(21) The instructions **216** can be one or more instructions that when executed by one or more processors of the computing platform **105**, causes the one or more processors to perform actions defined by the instructions. The instructions **216** can be stored in object code format for direct processing by the processors **212**, or other in other formats including interpretable scripts or collections of independent source code modules that are interpreted on demand or compiled in advance.

(22) The data **218** can be retrieved, stored or modified by the processors **212** in accordance with the instructions **216**. The data **218** can be stored in computer registers, in a relational or non-relational database as a table having a plurality of different fields and records, or as JSON, YAML, proto, or XML documents. The data **218** can also be formatted in a computer-readable format such as, but not limited to, binary values, ASCII or Unicode. Moreover, the data **218** can include information sufficient to identify relevant information, such as numbers, descriptive text, proprietary codes, pointers, references to data stored in other memories, including other network locations, or information that is used by a function to calculate relevant data.

(23) The network computing device **110** also includes a network protocol stack **305**. The network protocol stack **305** can include different software and/or hardware layers for communicating across the network **125** according to a particular communication protocol. Different layers can communicate with different components of the network computing device **110**. At the highest layer, the network protocol stack **305** can include a layer for communicating with user applications. At lower layers, the network protocol stack **305** is configured to process data from the higher layers to communicate the data to nodes of the network **125**. As part of the network protocol stack **305**, the network protocol stack **305** passes data packets to the NIC **220** to be transmitted to a destination node. Example protocols that the network protocol stack **305** can implement include HTTP, TCP, IP, and Ethernet. Other example protocols include the user datagram protocol (UDP), the reliable datagram protocol (RDP), the reliable user datagram protocol (RUDP), and the datagram congestion control protocol (DCCP).

(24) A data packet can be a collection of data of a fixed size, for example between 20 to 512 or more bytes, which follows a specific ordering of data and/or metadata in the data packet. The exact size of a data packet can vary depending on the configuration of the network, the nodes sending or receiving the data packet, and/or the communication protocol by which the network **125** communicates the data packet. For example, a data packet configured for transmission over the transmission control protocol (TCP) can include one or more bits dedicated to a header, identification of a source and/or destination for the data packet, a checksum for verifying the contents of the data packet, and the payload or content to be transmitted. Data packets can include more or fewer fields from implementation-to-implementation.

(25) In some implementations in which the network computing device **110** is a physical computer, the network protocol stack **305** can be part of the operating system running on the network computing device **110**. In some implementations in which the network computing device **110** runs one or more virtual machines, the network protocol stack **305** can refer to a respective network protocol stack of each virtual machine, and a network protocol stack of the operating system running the one or more virtual machines. The network computing device **110** can include a hypervisor or some other mechanism for receiving packets from the one or more virtual machines and passing the packets to the network protocol stack of the operating system running the virtual machine(s).

(26) The packets can be generated by applications or services running on the network computing device **110**, directly on the operating system of the network computing device **110** and/or through one or more virtual machines running on the network computing device **110**.

(27) The network nodes **112A-N** can also be configured similar to the network computing device **110**. For example, the network node A **112A** can include one or more processors **213**, memory **215**, instructions **217**, data **219**, and a NIC **240**. In some implementations, the NIC **240** implements the traffic shaping engine **100** according to aspects of the disclosure described herein.

(28) The network computing device **110** can also include a user output **225**. The network computing device **110** can be configured to display at least a portion of received data on a display implemented as part of the user output **225**. The user output **225** can also be used for displaying an interface between the network computing device **110** and the computing platform **105**. The user output **225** can alternatively or additionally include one or more speakers, transducers or other audio outputs, a haptic interface or other tactile feedback that provides non-visual and non-audible information to a user of the network computing device **110**.

(29) Although FIG. 2 illustrates the processors **212**, **213** and the memories **214**, **215** as being within the computing devices **115**, **120**, components described in this specification, including the processors **212**, **213** and the memories **214**, **215** can include multiple processors and memories that can operate in different physical locations and not within the same computing device or node. For example, some of the instructions **216**, **217** and the data **218**, **219** can be stored on a removable CD-ROM or another removable storage medium, such as a read-only computer chip. Some or all of the instructions and data can be stored in a location physically remote from, yet still accessible by, the processors **212**, **213**. Similarly, the processors **212**, **213** can include a collection of processors that can perform concurrent and/or sequential operation. The network computing device **110** and the network nodes **112A-N** can each include one or more internal clocks providing timing information, which can be used for time measurement for operations and programs run by the network computing device **110** and the network nodes **112A-N**.

(30) FIG. 3 is a relational block diagram illustrating the traffic shaping engine **100** traffic shaping outbound data packets. The traffic shaping engine includes a time-indexed data structure **301** and a scheduler **310**. The traffic shaping engine **100** receives data packets **303** and transmits the packets to one or more destination nodes on the network **125**.

(31) In some implementations, the one or more destination nodes and/or other nodes communicating data packets to the network computing device **110** can also implement the traffic shaping engine **100**. Aspects of the disclosure described herein with respect to the network computing device **110** can also be implemented on at least one of the nodes in communication with the network computing device **110**.

(32) The traffic shaping engine **100** can transmit the data packets **303** without delay. The traffic shaping engine **100** processes the data packets as they are received, and generates confirmation tokens **306**, one token per data packet. A confirmation token is data reported from the NIC **220** to the network protocol stack **305** to signal to the network protocol stack **305** to push another data packet to the NIC **220** for transmission.

(33) In conventional controllers, a confirmation token can be generated when a corresponding data packet has been transmitted. However, as described according to aspects of this disclosure herein, a confirmation token can be sent before or after a corresponding packet for the confirmation token is transmitted. In other words, the generation of a confirmation token is independent of whether or not the corresponding data packet was transmitted to the network **125** by the NIC **220**. The NIC **220**, through the traffic shaping engine **100**, can manage the rate at which additional packets are pushed by the network protocol stack **305** to the traffic shaping engine **100**, by selectively dequeuing confirmation tokens stored by the time-indexed data structure **301** and the scheduler **310**.

(34) The time-indexed data structure **301** can include one or more data structures configured to store and fetch data elements based on a timestamp associated with the element. The timestamp can

be relative to a standard time, such as an internal clock maintained by the network computing device **110**. The timestamp can also be relative to other timestamps, for example different elements have ordinal timestamps 1.sup.st, 2.sup.nd, 3.sup.rd, etc., to indicate each element's position in time relative to one another. A first-in first-out (FIFO) queue can be an example of a time-indexed data structure, because data is stored according to the time at which the data entered the queue. Other types of time-indexed structures include a timing-wheel and a calendar queue. As another example, time-indexed data structures can include any of a variety of other types of priority queues, including first-in last-out (FILO) or last-in first-out (LIFO).

(35) The time-indexed data structure **301** can be configured to include multiple positions to store elements up to a maximum period of time into the future. For example, a time-indexed data structure can include 50 positions, each position representing a corresponding difference in time between elements stored in adjacent positions. If each position represented, for example, 2 milliseconds, then the time-indexed data structure as a whole can represent elements up to 100 milliseconds into the future. In this example, the position at index 0 would store an element corresponding to the current time, and the position at index 49 would store an element corresponding to 100 milliseconds ahead of the current time.

(36) Another example of a time-indexed data structure is a timing wheel. A timing wheel can be implemented as a queue, with each element in the queue corresponding to a certain offset relative to a current time. The time-indexed data structure **301** can include a timing wheel, and in some implementations, additional data structures for storing elements, such as confirmation tokens as described herein.

(37) As an example, the additional data structures that can form part of the time-indexed data structure can include socket buffers. Socket buffers store data sent from sockets used by different applications running on the network computing device **110** for transmitting data to the network interface controller **220**. A socket can be a software component for communicating data from an application to a lower layer of the network protocol stack **305**. The socket can be programmatically defined, for example according to an application program interface (API), and an application running on the network computing device **110** can generate and destroy one or more sockets as appropriate for transmitting data to the network interface controller **220**. The socket buffers of the time-indexed data structure **301** can store confirmation tokens corresponding to packets generated and sent to the traffic shaping engine **100** by various applications running on the network computing device **110**.

(38) The time-indexed data structure **301** is configured to enqueue the confirmation tokens **306** for the data packets **303**. Confirmation tokens can be generated as the traffic shaping engine **100** receives the data packets **303**, however as described herein, the enqueueing of a confirmation token is independent of whether the corresponding packet for the confirmation token has been sent or not. In other words, each of the confirmation tokens **306** may be enqueued in the time-indexed data structure **301** before or after each corresponding data packet of the data packets **303** is transmitted to the network **125**.

(39) The scheduler **310** can timestamp each of the confirmation tokens **306** according to a traffic shaping policy. The scheduler **310** can control the rate at which dequeued confirmation tokens **308** are sent to the network protocol stack **305**, thereby controlling the rate at which additional packets are pushed to the NIC **220** by the network protocol stack **305**. The network protocol stack **305** can be configured to transmit additional packets to the network interface controller **220** in response to receiving confirmation tokens.

(40) In some implementations, the scheduler **310** does not explicitly timestamp each of the confirmation tokens **306**, but instead manages the order in which they are enqueued sequentially. The time-indexed data structure **301** can then dequeue or pop confirmation tokens at a predetermined rate, and the order at which the time-indexed data structure **301** dequeues the confirmation tokens can be in the order the scheduler **310** scheduled the confirmation tokens, in

accordance with the traffic shaping policy.

(41) From the perspective of the network protocol stack **305**, the stack **305** assumes that the corresponding packet for a received dequeued confirmation token has been sent to the network **125**. In practice, however, this may or may not be true as of the time the network protocol stack **305** receives the dequeued confirmation token.

(42) For example, the dequeued confirmation token may correspond to a packet that has not yet been transmitted by the traffic shaping engine **100**. However, at least because the scheduler **310** can timestamp confirmation tokens for different future times, the scheduler **310** can cause a confirmation token to be dequeued and sent to the network protocol stack **305** such that a corresponding packet for the confirmation token has likely transmitted by the time the network protocol stack **305** pushes another packet. The timing by the scheduler **310** coupled with the traffic shaping engine **100** not delaying the packets **303** can create a dataflow in which confirmation tokens are not generated strictly on completion of the transmission of the corresponding packets, but are sent to the network protocol stack **310** such that the network protocol stack **305** is not affected by actual discrepancies between the confirmation token and the actual transmission.

(43) Instead of delaying packets and storing packets in memory coupled to the network interface controller **220**, the network interface controller **220** can instead store the confirmation tokens for the packets, and transmit the packets themselves without delay. Confirmation tokens can be much smaller relative to a packet, for example 16 bytes for a confirmation token versus 512 bytes for a corresponding packet.

(44) Network interface controllers implemented according to aspects of this disclosure can be implemented with less memory overall, and the smaller amounts of memory that is used can be of a higher quality that is cost-prohibitive for storing entire data packets. Further, network interface controllers that are bottlenecked by memory capacity can operate faster, at least because processing cycles are not wasted waiting for memory to become available for storing additional data packets.

(45) Aspects of this disclosure can also provide for advantages over approaches in which memory pointers are stored instead of a corresponding packet waiting transmission. This is at least because the network interface controller **220** does not require storing memory pointers for data packets which can be between 4-8 bytes for each delayed packet. Further, there is a fixed computational cost to store the memory pointer, determine that the data packet is ready for transmission, load the memory pointer, and fetch the data packet at the address in memory indicated by the pointer. The network interface controller **220**, by contrast, can transmit the packet without these additional steps, all while saving on memory usage and still implementing a traffic shaping policy to prevent or mitigate network congestion or blocking.

(46) The scheduler **310** can be configured to timestamp confirmation tokens before the elements are enqueued into the time-indexed data structure **301**. The timestamp of an element stored in a time-indexed data structure can represent the time at which the element should be dequeued. The timestamp can be relative to a current time, for example 10 milliseconds in the future from the current time. Multiple elements can be timestamped with the same time. In those cases, the time-indexed data structure **301** can dequeue the same-timestamped elements according to an additional strategy, such as first-in-first-out dequeuing.

(47) As described herein, the scheduler **310** can timestamp confirmation tokens according to a traffic shaping policy. A traffic shaping policy can refer to one or more heuristics, rules, or algorithms for the scheduler **310** to determine how incoming confirmation tokens are timestamped. Generally, the scheduler **310** timestamps higher priority confirmation tokens with a timestamp closer to the current time, and timestamps lower priority confirmation tokens with a timestamp farther from the current time.

(48) The scheduler **310** can update timestamps for confirmation tokens in response to confirmation tokens received later-in-time. For example, the scheduler **310** may receive one or more confirmation tokens for packets originating from a higher priority application executing on the

network computing device **110** as compared with currently enqueued tokens. In response, the scheduler **310** can update the timestamps of the previously enqueued confirmation tokens to prioritize the newly received confirmation tokens.

(49) The traffic shaping policy can be one policy or a composite of different policies, each of the different policies corresponding to a different source of packets from the network protocol stack **305**. For example, the network computing device **110** can execute a plurality of different applications that emit data packets, and the scheduler **310** can schedule data packets from each application in accordance with a respective traffic shaping policy. The traffic shaping policy can specify how different data packets should be prioritized. For example, the traffic shaping policy can specify that data packets from a particular source, such as a particular application running on the network computing device **110**, be prioritized over data packets from other applications.

(50) Different policies can relate to other criteria besides packet source. One example can be the type of packet being transmitted, such as different data packets for different communication protocols. The scheduler **310** can store and access different policies according to some data structure mapping an incoming confirmation token to a policy associated with the corresponding data packet for that token. One example of such a data structure can be a hash table.

(51) The traffic shaping policy can also specify how the scheduler **310** should rate-limit and/or pace incoming traffic. Pacing can refer to injecting inter-packet gaps to smooth traffic, which can be particularly useful when the network interface controller **220** interfaces with the network **125** over a single connection. Rate-limiting can refer to the rate at which data packets are transmitted, which can be controlled by the traffic shaping engine **100** through the scheduler **310**. As an example, the scheduler **310** can schedule dequeued confirmation tokens to receive additional data packets from the network protocol stack **305** to a rate at which the network interface controller **220** can sustain to continuously transmit data packets without delay. The rate can be based on the physical limitations of the hardware the traffic shaping engine **100** is implemented on, for example because the hardware can only transmit so many data packets in a second.

(52) The traffic shaping policy can specify transmitter-side and receiver-side rate limiting. In other words, the scheduler **310** can schedule confirmation tokens to be dequeued based on data packets received from the network protocol stack **305** and/or data packets received from other nodes on the network **125**. One or more nodes transmitting data packets to the network computing device **110** can also implement a traffic shaping engine similar to the traffic shaping engine **100**. The transmitting node(s) and other nodes coupled to the network **125** can be configured to implement respective traffic shaping policies.

(53) In some implementations, the scheduler **310** can also schedule confirmation tokens based on packet acknowledgments sent from one or more nodes that have received data packets transmitted by the network computing device **110**. A packet acknowledgment is data sent in response to a successful receipt of a data packet from a receiving network node of the network **125**.

(54) In some implementations, the network protocol stack **305** is configured to automatically push a data packet to the traffic shaping engine **100** in response to receiving a packet acknowledgment. In some implementations, the traffic shaping engine **100** receives the packet acknowledgment and can cause the scheduler **310** to schedule confirmation tokens for dequeuing at least partially based on receipt of packet acknowledgments. For example, receipt of packet acknowledgments from a particular network node can indicate that the network node is available to receive more data. In response, the scheduler **310** can prioritize dequeuing tokens corresponding to data packets for an application in communication with the recipient network node.

(55) The scheduler **310** can be configured to dequeue confirmation tokens in exchange for additional packets from the network protocol stack **305** so as to provide backpressure as necessary for the network interface controller **220**. Backpressure can refer to the opposition of additional incoming data to a component subject to one or more conditions. For example, the dequeued confirmation tokens **308** can provide a form of backpressure between the traffic shaping engine **100**

and the network protocol stack **305**. This is at least because the network protocol stack **305** can be configured to emit additional packets to the traffic shaping engine **100** only in response to the dequeued confirmation tokens **308**. Because the scheduler **310** determines when confirmation tokens are dequeued, the scheduler **310** can adjust backpressure between the network protocol stack **305** and the network interface controller **220** to reduce or increase the stream of incoming data packets as needed.

(56) FIG. **4** is a block diagram of the time-indexed data structure **301** for the traffic shaping engine **100**. As described herein with reference to FIG. **3**, the network protocol stack **305** can push the packets **303** to the traffic shaping engine **100**. The traffic shaping engine **100** can transmit the packets **303** and can generate the confirmation tokens **306**. The scheduler **310** can receive the confirmation tokens **310** and can schedule their entry into the time-indexed data structure **301**.

(57) The time-indexed data structure **301** as shown in FIG. **4** can include a socket buffer A **403A** and a socket buffer B **403B**. The socket buffers **403A-B** can correspond to confirmation tokens corresponding data packets generated by applications **402A-B**, respectively. T_A1, T_A2, and T_A3 refer to confirmation tokens for data packets A1, A2, and A3 generated by the application A **402A**. Similarly, T_B1, T_B2, and T_B3 refer to confirmation tokens for data packets B1, B2, and B3 generated by the application B **402B**. Consistent with aspects of this disclosure, the scheduler **310** can enqueue the confirmation tokens T_A1 through T_B3 independent of whether or not the traffic shaping engine **100** transmitted the corresponding data packets A1-B3.

(58) The time-indexed data structure **301** can also include a timing wheel **405**. The timing wheel **405** can be implemented as a circular queue, in which the timing wheel **405** corresponds to a future time period, and each position of the timing wheel **405** corresponds to an interval of time within the future time period.

(59) For ease of description, the timing wheel **405** is described as having positions with respective indices labeled 0-7. T_B1 is shown in the position at index 0, and T_A1 is shown in the position at index 3. Also for ease of description, the timing wheel **405** is shown as rotating clockwise through time. It is understood, however, that in some implementations the timing wheel **405** can include hundreds or thousands of positions with corresponding indices.

(60) As the timing wheel **405** moves forward in time, the scheduler **310** can check whether a position at the current time **401** has a confirmation token to dequeue. If so, the scheduler **310** can dequeue the confirmation token and send the dequeued confirmation token to the network protocol stack **305** in exchange for another data packet. The scheduler **310** can schedule confirmation tokens intended to be dequeued within the future time period specified by the timing wheel **405**. The scheduler **310** can buffer confirmation tokens that are scheduled later than the future time period in one of the socket buffers **403A-B**.

(61) The scheduler **310** can be configured to poll the time-indexed data structure **301** at a constant interval to determine whether a confirmation token is scheduled to be dequeued, for example because the timestamp is the same time or older than the current time. The scheduler **310** can poll the time-indexed data structure **301** can be polled periodically with a period equal to the length of time associated with each time slot or a multiple thereof.

(62) The traffic shaping policy implemented by the scheduler **310** can specify how confirmation tokens are dequeued from the socket buffers **403A-B**. The traffic shaping policy can specify rules for balancing the dequeuing of each socket buffer fairly and/or according to some weight to prefer dequeuing by one socket buffer over another. For example, the scheduler **310** can alternate between the socket buffers **403A-B**. As another example, the scheduler **310** can prioritize dequeuing confirmation tokens closer to the current time, and/or prioritize dequeuing confirmation tokens from the application A **402A** over the application B **402B**.

(63) The scheduler **310** can also schedule the enqueued confirmation tokens according to a rate limit. The scheduler **310** can timestamp confirmation tokens for later times to allow for some confirmation tokens to be dequeued later relative to others, which can help to maintain the current

rate at which tokens are dequeued and new packets are pushed to the network interface controller **220** and the traffic shaping engine **100**.

(64) FIG. 5 is a block diagram of a traffic shaping engine **500** implementing separate time-indexed data structures **501**, **506** for delayed and un-delayed outbound data packets. The traffic shaping engine **500** according to some aspects of this disclosure can transmit some data packets without delay according to a first traffic shaping policy, such as the traffic shaping policy described herein with reference to FIG. 3. The traffic shaping engine **500** can additionally delay some other data packets according to a second traffic shaping policy. The second traffic shaping policy can be implemented by a scheduler **504** and a time-indexed data structure **506**. In some implementations, both the first and the second traffic shaping policies can be implemented by the same scheduler in communication with the time-indexed data structures **501**, **506**.

(65) The traffic shaping engine **500** can receive packets **502** from the network protocol stack **305** that are intercepted by the scheduler **510**. The scheduler **510** can be configured to flag some packets **503** for immediate transmission, without delay. The scheduler **510** can flag packets according to a traffic shaping policy. For example, the traffic shaping policy may specify flagging packets that are subsequent to a delayed packet in time. The scheduler **510** can query the memory **505** to determine if the incoming packets follow any data packets that are currently delayed. By flagging some packets for un-delayed transmission, the traffic shaping engine **500** can mitigate or eliminate head-of-line blocking. Head-of-line blocking can occur when packets later-in-time are held up by a packet blocked that is earlier-in-time.

(66) One way to mitigate or eliminate head-of-line blocking is to cause data packets received by the network interface controller to schedule data packets out of order. When the transmission of a data packet is blocked, for example because the recipient network node is currently unable to receive additional data packets, the traffic shaping engine **500** can flag one or more data packets received later-in-time for un-delayed transmission. In general, the traffic shaping engine **500** can flag data packets whose transmission is unrelated to the transmission of the blocked data packet. For example, the flagged data packets can include data packets that are not in communication with the blocked receipt network node.

(67) However data packets are flagged, the traffic shaping engine **500** can transmit the flagged packets **503** without delay, in accordance with aspects of the disclosure described herein with reference to FIGS. 3-4. Other data packets that are not flagged may pass through the traffic shaping engine **500** without delay, or may be delayed for a variety of reasons. For example, data packets may be delayed because their transmission is blocked by a recipient network node, or because the rate limit for the network interface controller implementing the traffic shaping engine **100** has been exceeded. The traffic shaping engine **500** is configured to prioritize the transmission of the flagged packets **503** before other data packets. As an example, if the traffic shaping engine **500** flags a data packet for un-delayed transmission, the network interface controller implementing the traffic shaping engine **500** can transmit the flagged data packet before any un-flagged data packet scheduled for transmission.

(68) The scheduler **504** can schedule the transmission of delayed packets **505**. The time-indexed data structure **506** can be of any type of time-indexed data structure as described herein with reference to the time-indexed data structure **301**, with one difference being that the time-indexed data structure **506** is configured to enqueue packet identifiers **507** corresponding to the delayed packets **505**. A packet identifier is data associated with a delayed data packet stored in memory **505** coupled to the traffic shaping engine **500**. A delayed data packet can be stored in the memory **505** until the time the corresponding packet identifier for the data packet is scheduled to be dequeued from the time-indexed data structure **506**.

(69) The scheduler **504** can timestamp the packet identifiers **507** according to the second traffic shaping policy and in accordance with description in this disclosure herein for the scheduler **310** in FIG. 3. The second traffic shaping policy can be any traffic shaping policy that specifies rules,

heuristics, algorithms, or other conditions for delaying the transmission data from a network interface controller. The traffic shaping engine **500** illustrates how aspects of the disclosure can be implemented to augment existing solutions to traffic shaping in a discrete manner. Rather than modify the network protocol stack or other components of a network computing device, additional logic can be added to an existing network interface controller that can improve data transmission at least because of the mitigation or elimination of head-of-line blocking in accordance with aspects of this disclosure.

(70) In some implementations in which the network interface controller is programmable, the network interface controller can be programmed consistent with aspects of this disclosure with minimal to no changes to other components of the network computing device. As a result, traffic shaping engines such as the traffic shaping engines **100**, **500** can be implemented on a large scale to hundreds or thousands of network nodes in a manner that is agnostic to the communication protocol deployed by a network and its network architecture.

(71) FIG. **6** is a flowchart of an example process **600** for traffic shaping on a NIC. The process **600** is described as being performed by a network computing device implementing a traffic shaping engine and programmed in accordance with aspects of this disclosure. For example, the network computing device **110** implementing the traffic shaping engine **100**, appropriately programmed, can perform the process **600**. Steps of the process **600** can be performed in different orders, with additional steps added or some steps as shown in the process **600** removed.

(72) The network computing device receives **605** data packets for transmission to one or more network nodes of a network and according to a traffic shaping policy. As described herein with reference to FIGS. **3-5**, the network computing device can receive the data packets from a network protocol stack implemented on the network computing device. The packets can be received by a network interface controller implementing the traffic shaping engine.

(73) The network computing device generates and stores **610** a first confirmation token for the first data packet in a time-indexed data structure. As described herein with reference to FIGS. **3-5**, the traffic shaping engine can implement a time-indexed data structure and a scheduler that is configured to timestamp and insert confirmation tokens according to a timestamp. The scheduler uses the traffic shaping policy to determine what timestamp is given to each confirmation token.

(74) The network computing device transmits **615** the first data packet to the one or more network nodes. As described herein with reference to FIGS. **3-5**, the network computing device can transmit the first data packet without delay to one or more network nodes of the network. The transmission of the first data packet can be independent of the generation and enqueueing of the first confirmation token corresponding to the first data packet.

(75) The network computing device dequeues **620** the first confirmation token for the first data packet according to the traffic shaping policy. The scheduler of the traffic shaping engine can be configured to timestamp the first confirmation token at a time relative to a current time but within a future time period, such as 100 milliseconds. As the current time increments, the scheduler can poll the time-indexed data structure to determine if there is a confirmation token to dequeue at the current time. The scheduler can dequeue the first confirmation token when the timestamp for the first confirmation token reaches the current time. The traffic shaping engine **100** can set the future time period in accordance with a rate limit defined in the traffic shaping policy. For example, to enforce a rate of 1 megabit of data per second, the future time period might extend to 12 milliseconds from the current time.

(76) After the network computing device dequeues **620** the first confirmation token, the network computing device can repeat the process **600** and receive **605** data packets for transmission. The received data packets can be, for example, data packets received in response to dequeuing **620** the first confirmation token.

(77) Aspects of this disclosure can provide for the following technical advantages. A network interface controller (NIC) can be configured to delay, pace, or rate packets of data while avoiding

transmission delay. Rather than saving packets to memory to be shaped during transmission, a NIC can pass at least some packets without delay for transmission and limit additional packets to the NIC according to enqueued confirmation tokens. A confirmation token is generated for each data packet processed through the NIC and is generally smaller in size than a data packet, for example 8 bytes versus 512 bytes or more of a data packet.

(78) Confirmation tokens can be enqueued and dequeued according to a traffic shaping policy, independent of the corresponding data packets that are transmitted without delay. Despite this independent processing of confirmation token and packet, a NIC implemented in accordance with aspects of this disclosure can effectively shape traffic while avoiding delay of transmission of the data packets: bottlenecking at the network computing device implementing the NIC, bursty traffic patterns, and/or head-of-line blocking caused by blocked packets earlier-in-time holding up subsequent packets.

(79) A programmable NIC can be programmed in accordance with aspects of this disclosure without modification to the network protocol stack of a host network device. Traffic shaping can be performed locally to distribute the computational burden of traffic shaping to intermediate devices dedicated between nodes of a network. For at least this reason, aspects of this disclosure provide for an easily scalable approach to traffic shaping, even for networks with hundreds or thousands of inter-communicating nodes running a variety of different applications and/or services.

(80) A NIC can be implemented in accordance with aspects of this disclosure to discriminate and expedite transmission of different data packets according to various traffic shaping policies. For example, some packets may be flagged by the NIC for expedited transmission. Flagged packets can be transmitted without delay and their confirmation tokens can be enqueued and dequeued according to a time-indexed data structure. For un-flagged packets, the NIC can implement any of a variety of other traffic shaping techniques, including techniques that delay transmission of some packets and store delayed packets to memory coupled to the NIC. Because the flagged packets can be transmitted without delay, a NIC implemented in accordance with aspects of this disclosure can use less memory overall than a NIC implementing conventional traffic shaping techniques, and can reduce or eliminate delays caused by head-of-line blocking.

(81) Aspects of this disclosure can be implemented in digital circuits, computer-readable storage media, as one or more computer programs, or a combination of one or more of the foregoing. The computer-readable storage media can be non-transitory, for example as one or more instructions executable by a cloud computing platform and stored on a tangible storage device.

(82) A computer program can be written in any type of programming language, and according to any programming paradigm, such as declarative, procedural, assembly, object-oriented, data-oriented, functional, or imperative. A computer program can be written to perform one or more different functions and to operate within a computing environment, such as on a physical device, virtual machine, or across multiple devices. A computer program can also implement functionality described in this specification as performed by a system, engine, module, or model.

(83) In this specification the phrase “configured to” is used in different contexts related to computer systems, hardware, or part of a computer program, module, or engine. When a system is said to be configured to perform one or more operations, this means that the system has appropriate software, firmware, and/or hardware installed on the system that, when in operation, causes the system to perform the one or more operations. When some hardware is said to be configured to perform one or more operations, this means that the hardware includes one or more circuits that, when in operation, receive input and generate output according to the input and corresponding to the one or more operations. When a computer program, module, or engine, is said to be configured to perform one or more operations, this means that the computer program includes one or more program instructions, that when executed by one or more computers, causes the one or more computers to perform the one or more operations.

(84) Unless otherwise stated, the foregoing alternative examples are not mutually exclusive, but

may be implemented in various combinations to achieve unique advantages. As these and other variations and combinations of the features discussed above can be utilized without departing from the subject matter defined by the claims, the foregoing description of the implementations should be taken by way of illustration rather than by way of limitation of the subject matter defined by the claims. In addition, the provision of the examples described herein, as well as clauses phrased as “such as,” “including” and the like, should not be interpreted as limiting the subject matter of the claims to the specific examples; rather, the examples are intended to illustrate only one of many possible embodiments. Further, the same reference numbers in different drawings can identify the same or similar elements.

Claims

1. A network computing device comprising one or more processors configured to: receive a first data packet for transmission to one or more network nodes of a network according to a first traffic shaping policy; transmit the first data packet to the one or more network nodes; enqueue a first confirmation token for the first data packet by storing the first confirmation token in a first time-indexed data structure at a time after transmitting the first data packet, the first confirmation token stored in association with a timestamp representing a time the first confirmation token is scheduled to be dequeued based on a rate limit for receiving data packets; dequeue the first confirmation token for the first data packet according to the first traffic shaping policy based on the timestamp; receive a second data packet for transmission to the one or more network nodes in response to dequeuing the first confirmation token; store the second data packet; enqueue an identifier for the second data packet by storing the identifier in a second time-indexed data structure; dequeue the identifier for the second data packet according to a second traffic shaping policy; and transmit the second data packet to the one or more network nodes in response to dequeuing the identifier.
2. The network computing device of claim 1, wherein the one or more processors are further configured to flag the first data packet for immediate transmission to the one or more network nodes.
3. The network computing device of claim 1, wherein the one or more processors are further configured to immediately transmit the first data packet to the one or more network nodes.
4. The network computing device of claim 1, wherein the one or more processors are further configured to enqueue a second confirmation token for the second data packet before transmitting the second data packet.
5. The network computing device of claim 1, wherein enqueueing the first confirmation token further comprises storing the first confirmation token in a socket buffer.
6. The network computing device of claim 1, wherein the one or more processors are further configured to control a rate at which confirmation tokens are dequeued.
7. The network computing device of claim 6, wherein the one or more processors are further configured to enqueue the first confirmation token to maintain the rate at which confirmation tokens are dequeued.
8. The network computing device of claim 1, wherein: the one or more processors are coupled to a network interface controller; and the one or more processors are further configured to transmit the first data packet to the one or more network nodes using the network interface controller.
9. A method comprising: receiving, by a network computing device, a first data packet for transmission to one or more network nodes of a network according to a first traffic shaping policy; transmitting, by the network computing device, the first data packet to the one or more network nodes; enqueueing, by the network computing device, a first confirmation token for the first data packet by storing the first confirmation token in a first time-indexed data structure at a time after transmitting the first data packet, the first confirmation token stored in association with a timestamp representing a time the first confirmation token is scheduled to be dequeued based on a

rate limit for receiving data packets; dequeuing, by the network computing device, the first confirmation token for the first data packet according to the first traffic shaping policy based on the timestamp; receiving, by the network computing device, a second data packet for transmission to the one or more network nodes in response to dequeuing the first confirmation token; storing, by the network computing device, the second data packet; enqueueing, by the network computing device, an identifier for the second data packet by storing the identifier in a second time-indexed data structure; dequeuing, by the network computing device, the identifier for the second data packet according to a second traffic shaping policy; and transmitting, by the network computing device, the second data packet to the one or more network nodes in response to dequeuing the identifier.

10. The method of claim 9, further comprising flagging, by the network computing device, the first data packet for immediate transmission to the one or more network nodes.

11. The method of claim 9, further comprising immediately transmitting, by the network computing device, the first data packet to the one or more network nodes.

12. The method of claim 9, further comprising enqueueing, by the network computing device, a second confirmation token for the second data packet before transmitting the second data packet.

13. The method of claim 9, wherein enqueueing the first confirmation token further comprises storing the first confirmation token in a socket buffer.

14. The method of claim 9, further comprising controlling, by the network computing device, a rate at which confirmation tokens are dequeued.

15. The method of claim 14, further comprising enqueueing, by the network computing device, the first confirmation token to maintain the rate at which confirmation tokens are dequeued.

16. A non-transitory computer-readable medium for storing instructions that, when executed by one or more processors, cause the one or more processors to perform a method comprising: receiving a first data packet for transmission to one or more network nodes of a network according to a first traffic shaping policy; transmitting the first data packet to the one or more network nodes; enqueueing a first confirmation token for the first data packet by storing the first confirmation token in a first time-indexed data structure at a time after transmitting the first data packet, the first confirmation token stored in association with a timestamp representing a time the first confirmation token is scheduled to be dequeued based on a rate limit for receiving data packets; dequeuing the first confirmation token for the first data packet according to the first traffic shaping policy based on the timestamp; receiving a second data packet for transmission to the one or more network nodes in response to dequeuing the first confirmation token; storing the second data packet; enqueueing an identifier for the second data packet by storing the identifier in a second time-indexed data structure; dequeuing the identifier for the second data packet according to a second traffic shaping policy; and transmitting the second data packet to the one or more network nodes in response to dequeuing the identifier.
