

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12418511

Kind Code

B2

Date of Patent

September 16, 2025

Inventor(s)

Hohne; Charles C. et al.

Method and apparatus for secure communication and routing

Abstract

An apparatus is provided, comprising: a volatile memory; a non-volatile memory; a first electronic circuit that is configured to operate as a wireless access point, the first electronic circuit including a wireless controller for accessing a wireless network; and a second electronic circuit that is operatively coupled to the first electronic circuit, the second electronic circuit including at least one processor configured to execute: (i) a first virtual machine that includes a wireless network authentication server, and (ii) a second virtual machine that includes a virtual private network (VPN) server, wherein the wireless network authentication server is configured to authenticate devices that attempt to join the wireless network; wherein the VPN server is arranged to encrypt data that is received at the apparatus to produce encrypted data, and forward the encrypted data to the wireless controller for transmission over the wireless network.

Inventors: Hohne; Charles C. (Carmel, IN), Swails; Todd A. (Greenwood, IN), Nord; Christopher J. (Indianapolis, IN), Pulling; Christopher J. (Fishers, IN), Irgens; Kristi (Breckenridge, CO), Turner; Howard (McCordsville, IN), Knopf; Ian A. (Indianapolis, IN), Maglio; Vincent A. (Fishers, IN), Brown; Kyle A. (Fishers, IN)

Applicant: Vertex Aerospace, LLC (Madison, MS)

Family ID: 1000008821818

Assignee: Vertex Aerospace LLC (Madison, MS)

Appl. No.: 18/064332

Filed: December 12, 2022

Prior Publication Data

Document Identifier

Publication Date

US 20230111416 A1

Apr. 13, 2023

Related U.S. Application Data

Publication Classification

Int. Cl.: H04L9/40 (20220101); G06F9/455 (20180101); H04L12/46 (20060101); H04W12/037 (20210101); H04W12/06 (20210101); H04W12/088 (20210101); H04W12/121 (20210101)

U.S. Cl.:

CPC H04L63/0272 (20130101); G06F9/45558 (20130101); H04L12/4625 (20130101); H04L12/4641 (20130101); H04W12/037 (20210101); H04W12/06 (20130101); H04W12/088 (20210101); H04W12/121 (20210101); G06F2009/45583 (20130101); G06F2009/45595 (20130101)

Field of Classification Search

CPC: H04L (63/0272); H04W (12/06); H04W (12/121); H04W (12/037); H04W (12/088); G06F (9/45558); G06F (2009/45583); G06F (2009/45595)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2007/0189308	12/2006	Tchigevsky	370/396	H04L 12/4625
2020/0310972	12/2019	Shanbhogue	N/A	H04L 9/30
2021/0344651	12/2020	Joshi	N/A	H04L 12/4641

Primary Examiner: Anyan; Barbara B

Attorney, Agent or Firm: Jones Day

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a divisional application of U.S. patent application Ser. No. 16/910,898, filed Jun. 24, 2020, which is incorporated herein by reference in its entirety.

BACKGROUND

(1) A wireless network is a type of telecommunications network that uses radio communications to transmit data. Wireless networks are frequently connected to wired networks and used to access information stored in the wired networks. Wireless networks are inherently more vulnerable to attacks than wired networks, as they are accessible to anyone within range. This property of wireless networks makes them suitable staging grounds for attacks on connected wired networks. The provision of adequate security mechanisms in wireless networks is thus necessary to ensure safe and reliable operation of the wireless networks, as well as any connected wired networks.

SUMMARY

(2) According to aspects of the disclosure, an apparatus is provided, comprising: a volatile memory;

a non-volatile memory; a first electronic circuit that is configured to operate as a wireless access point, the first electronic circuit including a wireless controller for accessing a wireless network; and a second electronic circuit that is operatively coupled to the first electronic circuit, the second electronic circuit including at least one processor configured to execute: (i) a first virtual machine that includes a wireless network authentication server, and (ii) a second virtual machine that includes a virtual private network (VPN) server, wherein the wireless network authentication server is configured to authenticate devices that attempt to join the wireless network; wherein the VPN server is arranged to encrypt data that is received at the apparatus to produce encrypted data, and forward the encrypted data to the wireless controller for transmission over the wireless network, and wherein at least one of the first virtual machine or the second virtual machine is fully contained in the volatile memory.

(3) According to aspects of the disclosure, an apparatus is provided, comprising: a non-volatile memory; a volatile memory; a first electronic circuit that is that is configured to operate as a wireless access point, the first electronic circuit including a wireless controller for accessing a wireless network; and a second electronic circuit that is operatively coupled to the first electronic circuit, the second electronic circuit including at least one processor configured to execute: (i) a first virtual machine that includes a virtual private network (VPN) server, and (ii) a second virtual machine that includes a first firewall, wherein the VPN server is configured to: encrypt data that is received at the apparatus to produce encrypted data, and forward the encrypted data to the wireless controller for transmission over the wireless network, and wherein each of the first virtual machine and the second virtual machine is fully contained in the volatile memory.

(4) According to aspects of the disclosure, a method is provided for use in an electronic device that includes a first electronic circuit configured to operate as an access point for accessing a wireless network and a second electronic circuit having a volatile memory and at least one processor, the method comprising: instantiating a random-access memory (RAM) disk in the volatile memory of the second electronic circuit; partitioning the RAM disk into a plurality of partitions; launching a first virtual machine on the second electronic circuit, the first virtual machine being launched in a first partition of the RAM disk, the first virtual machine including a wireless network authentication server that is configured to authenticate devices that attempt to join the wireless network via the first electronic circuit; and launching a second virtual machine on the second electronic circuit, the second virtual machine being launched in a second partition of the RAM disk, the second virtual machine including a virtual private network (VPN) server that is configured to encrypt data that is received at the apparatus to produce encrypted data, and forward the encrypted data to the first electronic circuit for transmission over the wireless network, wherein the first virtual machine and the second virtual machine are fully contained in the volatile memory of the second electronic circuit.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The foregoing features may be more fully understood from the following description of the drawings in which:
- (2) FIG. 1 is a schematic diagram of an example of a communications system, according to aspects of the disclosure;
- (3) FIG. 2 is a schematic diagram of an example of a communications system, according to aspects of the disclosure;
- (4) FIG. 3 is a schematic diagram of an example of a gateway router (GR), according to aspects of the disclosure;
- (5) FIG. 4 is a schematic diagram of an example of a primary system-on-a module (SOM) that is

part of the GR of FIG. 3, according to aspects of the disclosure;

(6) FIG. 5 is flowchart of an example of a process that is performed by the GR of FIG. 3, according to aspects of the disclosure.

(7) FIG. 6 is a state diagram illustrating a set of operational states of the GR of FIG. 3, according to aspects of the disclosure; and

(8) FIG. 7 is a perspective side view of the GR of FIG. 3, according to aspects of the disclosure.

DETAILED DESCRIPTION

(9) FIG. 1 is a diagram of an example of a communications system **100**, according to aspects of the disclosure. The communications system **100** may be deployed aboard an aircraft and/or ground vehicles, and it can be used for the exchange of data between the pilots, personnel on board the aircraft, and ground personnel. The communications system **100** may include a gateway router (GR) **110**, a wireless network **120**, and a wired network **130**. The wireless network **120** may include an 802.11 network (e.g., a WiFi network), a Long-Term Evolution (LTE) network, a 5G network, and/or any other suitable type of network. The wired network **130** may include a local area network (LAN), a wide area network (WAN), an Ethernet network, an Infiniband network, the Internet, and/or any other suitable type of communications network. A plurality of tablets **122** may be connected to the wireless network. The wired network, on the other hand, may be connected to one or more multifunction displays (MFDs) **132**, a maintenance terminal **134**, a Miltope™ computing system **136**, and a navigation device **138**. In some implementations, the GR **110** may be used by the tablets **122** as an access point for connecting to the wireless network **120**. Additionally or alternatively, in some implementations, the GR **110** may be arranged to operate as a bridge between the wireless network **120** and the wired network **130**, thus allowing data to be exchanged between any of the devices **132-138** and any of the tablets **122**.

(10) FIG. 2 illustrates another example of a communications system where the GR **110** can be deployed. Shown in FIG. 2 is a communications system **200** including a wireless network **220** that is coupled to a wired network **230** via the GR **110**. The wireless network **220** may include an 802.11 network (e.g., a WiFi network), a Long-Term Evolution (LTE) network, a 5G network, and/or any other suitable type of network. The wired network **230** may include a local area network (LAN), a wide area network (WAN), an Ethernet network, an Infiniband network, the Internet, and/or any other suitable type of communications network. A plurality of wireless devices **222** may be connected to the wireless network **220**, and a plurality of wired devices **232** may be coupled to the wired network **230**. Each of the wireless devices **222** may include a smartphone, a tablet, a laptop, and/or any other device that includes an interface for connecting to the wireless network **220**. Each of the wired devices **232** may include a desktop computer, a laptop, a server, and/or any other suitable type of device that includes an interface for connecting to the wired network **230**. As noted above, the GR **110** may be arranged to operate as an access point for connecting to the wireless network **220** and/or a bridge between the wireless network **220** and the wired network **230**. The operation of the GR **110** is discussed further below with respect to FIGS. 3-7.

(11) FIG. 3 is a diagram of the GR **110**, according to aspects of the disclosure. The GR **110** may be coupled to a wireless network **302** and a wired network **304**. In operation, the GR **110** may be configured to operate as an access point for connecting client devices to the wireless network **302**, as well as a bridge for connecting the wireless network **302** to the wired network **304**. The GR **110** may include a primary system-on-a-module (SOM) **310** and a secondary SOM **320** that is coupled to the primary SOM **310**. The primary SOM **310** may be coupled to a switch **330** and one or more data connectors **340**. The secondary SOM **320** may be coupled to one or more antenna connectors **350**. In operation, the GR **110** may receive data, from the wired network **304**, on any of the data connectors **340** and route the received data to any of the antenna connectors **350** for transmission over the wireless network **302**. Similarly, the GR **110** may receive data, from the wireless network **302**, on any of the antenna connectors **350** and route the received data to any of the data connectors **340** for transmission over the wired network **304**.

(12) The primary SOM **310** may include a plurality of data ports **312**, a data port **313**, a data port **314**, a processor **316**, a volatile memory **318**, and a non-volatile memory **319**. According to the present example, each of the data ports **312**, **313**, and **314** is an Ethernet port. However, alternative implementations are possible in which any of the data ports **312**, **313**, and **314** is a different type of port, such as a USB port, or a parallel data port, etc. The processor **316** may include one or more of a general-purpose processor (e.g., an ARM-based processor, a RISC processor, an x86 processor, etc.), a special-purpose processor, an Application-Specific Integrated Circuit (ASIC), a Field-Programmable Gate Array (FPGA), and/or any other suitable type of processing circuitry. The volatile memory **318** may include any suitable type of volatile memory, such as synchronous dynamic random-access memory (SDRAM) for example. The non-volatile memory **319** may include any suitable type of non-volatile memory, such as electrically erasable programmable read-only memory (EEPROM), a hard disk (HD), a solid-state drive (SSD), or a non-volatile random-access memory (nvRAM) device, for example.

(13) The secondary SOM **320** may include a port **322**, a wireless controller **324**, and a plurality of access point (AP) radios **326**. According to the present example, the port **322** is an Ethernet port. However, alternative implementations are possible in which the port **322** includes another type of port, such as a USB port or a parallel data port, etc. The wireless controller **324** may include any suitable type of processing circuitry that is configured to transmit/receive data over a wireless network and/or manage access to the wireless network. By way of example, the wireless controller **324** may be configured to perform various access point functions, such as controlling the power and/or respective data rates of wireless channels that are available in the wireless network **302**. Additionally or alternatively, in some implementations, the wireless network controller may implement a data-link-layer switch for transmitting data over the wireless network **302**. The AP radios **326** may include one or more transceivers for transmitting data over a wireless network.

(14) The data connectors **340** may include any suitable type of electronic circuitry for transmitting and/or receiving data from the wired network **304**. In some implementations, the data connectors **340** may include one or more circular connectors **702** (shown in FIG. 7). However, it will be understood that the present disclosure is not limited to using any specific type connector. Although in the example, of FIG. 7 only two circular connectors **702** are depicted, it will be understood that each of the circular connectors may provide multiple Ethernet channels (or links). The antenna connectors **350** may include any suitable type of electronic circuitry for connecting the AP radios **326** to one or more antennas (not shown) that are arranged to transmit and receive data over the wireless network **302**. In some implementations, the antenna connectors **350** may include RP-SMA connectors **704** (shown in FIG. 7). However, it will be understood that the present disclosure is not limited to using any suitable type of antenna connector.

(15) The switch **330** may include a network-layer switch that is configured to multiplex some of the channels in the data connectors **340** onto the data port **313**. According to the present example, the switch **330** is a gigabit switch. However, it will be understood that the present disclosure is not limited to using any specific type of switch being used in the GR **110**. Additionally or alternatively, in some implementations, the switch **330** may be altogether omitted from the GR **110**.

(16) According to the example of FIG. 3, the secondary SOM **320** brings access point capabilities to the GR **110**, and it may be configured to operate as an access point to the wireless network **302**. As such, the secondary SOM **320** may be configured to receive data from the primary SOM **310** and transmit the received data over the wireless network **302**. Furthermore, the secondary SOM **320** may be configured to receive data from the wireless network **302** and forward the received data to primary SOM **310** for transmission over the wired network **304**. In some implementations, the secondary SOM **320** may encrypt data that is received from the primary SOM **310** before the data is transmitted over the wireless network **302**. Furthermore, the secondary SOM **320** may decrypt data that is received over the wireless network **302** before forwarding the received data to the primary SOM **310**. The encryption/decryption may be performed using AES and/or any other

encryption standard that can be used to secure the traffic in a wireless network.

(17) According to the example of FIG. 3, the primary SOM **310** brings data encryption, data monitoring, and data logging capabilities to the GR **110**. As is discussed further below with respect to FIG. 4, the primary SOM **310** may be configured to implement a virtual network including a firewall **412d** (shown in FIG. 4) and a Virtual Private Network (VPN) server **412c** (also shown in FIG. 4). The firewall **412d** may be configured to monitor and log data that is exchanged through the primary SOM **310** between the secondary SOM **320** and the wired network **304**. The VPN server **412c** may be configured to encrypt user data is provided by the primary SOM **310** to the secondary SOM **320** for transmission over the wireless network. This encryption may be separate from any additional encryption performed by the secondary SOM **320**. The VPN server **412c** may be further configured to decrypt user data that is received over the wireless network **302** before transmitting the user data over the wired network **304**. In some implementations, the VPN server **412c** may be arranged to establish a VPN channel between the GR **110** and at least one wireless device that is connected to the wireless network **302**.

(18) According to the example of FIG. 3, the primary SOM **310** brings device authentication capabilities to the secondary SOM **320**. Specifically, the primary SOM **310** may execute a wireless network authentication server **412e** (shown in FIG. 4). The server **412e** may be configured to authenticate any wireless devices that attempt to join the wireless network **302** via the secondary SOM **320** (and/or the GR **110**). In this regard, in addition to user data that is exchanged between the wireless network **302** and the wired network **304**, the primary SOM **310** and the secondary SOM **320** may also exchange communications that are used to validate the identity of devices (or users) that attempt to use the secondary SOM **320** (and/or GR **110**) as an access point for connecting to the wireless network **302**. In some implementations, the communications that are used to validate the identity of devices (or users) that attempt to connect to the wireless network **302** may be formatted in accordance with the Remote Authentication Dial-In User Service (RADIUS) protocol. However, it will be understood that the present disclosure is not limited to using any specific protocol for controlling access to the wireless network **302**.

(19) In some implementations, all communication between the primary SOM **310** and the secondary SOM **320** may be performed over an Ethernet connection that is established between data ports **314** and **322**. As can be readily appreciated, in some implementations, the data ports **314** and **322** may be used to transmit (both user data and communications for validating the identity of devices (or users) that attempt to join the wireless network **302**. In some implementations, the data ports **314** and **322** may be the only means for the exchange of data between the primary SOM **310** and the secondary SOM **320**. Limiting the number of communications paths between the primary SOM **310** and the secondary SOM **320** is advantageous because it reduces the pathways along which an intrusion into the wireless network **203** (or the secondary SOM **320**) could spread into the primary SOM **310** (and/or the wired network **304**).

(20) In some implementations, using different SOMs to implement the access point capabilities of the GR **110** and the wireless network device authentication capabilities of the GR **110** is advantageous because it could bolster the security of the GR **110**. Under the arrangement shown in FIG. 3, the wireless network authentication server **412e** is executed in non-shared memory with the secondary SOM **320**, thus making it less likely for any intrusion into the wireless network **302** (and/or secondary SOM **320**) to compromise the authentication mechanisms that are used by the wireless network **302**.

(21) FIG. 4 is a diagram of the primary SOM **310**, according to aspects of the disclosure. The primary SOM **310** (or processor **316**) may be configured to execute a hypervisor **400**. The hypervisor **400** may be configured to execute virtual machines **410e-f**, as shown.

(22) The virtual machine **410a** may be configured to execute a router application **412a** (herein router **412a**) and one or more user applications **414a**. The user application **414a** may include any suitable type of application, such as a video tracking application, or a medical application, for

example. In some implementations, the user application **414a** may be omitted from the virtual machine **410a**.

(23) The router **412a** may receive any packets that enter the primary SOM **310** through one of the data ports **312**. Afterwards, the router **412a** may route the received packets to one of: (A) the virtual machine **410B** (and/or firewall **412b**) and (B) the virtual machine **410c** (and/or VPN server **412c**). Specifically, the router **412a** may route to the virtual machine **410c** (and/or VPN server **412c**) all packets that are designated for devices in the wireless network **302**. On the other hand, the router **412a** may route to the virtual machine **410b** (and/or firewall **412b**) any packets that are designated for any of the virtual machines **410b-f** (and/or applications **412b-f**). In some implementations, any packet that is designated for a device in the wireless network **302** may include an address of the device in its destination field. In some implementations, any packet that is designated for one of the virtual machines **410b-f** may have an address of the virtual machine in its destination field. Packets that are designated for devices in the wireless network **302** are provided to the secondary SOM **320** for transmission over the wireless network **302**. Packets that are designated for any of the virtual machines **410B-F** are processed internally in the primary SOM **310**. Such packets may be used as a basis for performing software updates and/or configuration changes on any of the virtual machines **410b-f** and/or applications **412b-f**. Furthermore, such packets may be used for updating and/or changing one or more configuration settings of the hypervisor **400**.

(24) According to the present disclosure, performing configuration changes and/or software updates on any of the virtual machines **310b-f** and/or applications **412b-f** is referred to as “performing maintenance” on the virtual machines or applications. Similarly, performing configuration changes and/or software updates on the hypervisor **400** is referred to as “performing maintenance” on the hypervisor **400**. Any data that is used as a basis for performing maintenance on any of the hypervisor **400**, the virtual machines **410b-f** and/or applications **412b-f** is herein referred to as maintenance data. For example, the term “maintenance data” may refer to an instruction to perform a software update or change a configuration setting of one of the virtual machines **410a-f** and/or applications **412a-f**. As another example, the term “maintenance data” may refer to the payload of the instruction (e.g., a software update file or the new value of the configuration setting, etc.).

(25) The virtual machine **410b** may be configured to execute a firewall application **412b** (hereinafter firewall **412b**). The firewall **412b** may be configured to monitor data traffic (e.g., one or more data packets) that is designated for the virtual machine **410f** (and/or manager **412±**). As used throughout the disclosure, the phrase “monitoring data traffic” may refer to filtering the traffic, logging the traffic, and/or performing any other suitable action that is customarily performed by firewalls.

(26) The firewall **412b** may forward any traffic that is received at the firewall **412b** to a virtual switch **420** (and/or manager **412±**). The virtual switch **420** may be implemented as part of the hypervisor **400** or in a separate virtual machine (not shown). The virtual switch **420** may be configured to forward data traffic that is received from the firewall **412b** to the virtual machine **410f** (and/or manager **412±**). Furthermore, the virtual switch **420** may be configured to forward data traffic originating from the virtual machine **410f** (and/or manager **412±**) to any of the virtual machines **410b-e** and/or applications **410b-e**. Although in the present example, the virtual switch **420** is configured to route all traffic received from the firewall **412b** to the manager **412f**, alternative implementations are possible in which the virtual switch **420** routes traffic received from the firewall **412b** to any of the virtual machines **410b-f** and/or applications **412b-f**.

(27) The virtual machine **410f** may be configured to execute a manager application **412f** (hereinafter “manager **412f**”). In some implementations, the manager **412f** may be configured to receive maintenance data that is input into the GR **110**. The maintenance data may be received from a maintenance terminal, such as the maintenance terminal **134**, which is shown in FIG. 1. In some implementations, based on the maintenance data, the manager **412f** may perform a software update on any of the virtual machines **410b-f** and/or applications **412b-f**. Additionally or

alternatively, based on the maintenance data, the manager **412f** may change one or more configuration settings of any one of the virtual machines **410b-f** and/or applications **412b-f**. Additionally or alternatively, based on the maintenance data, the manager **412f** may change one or more configuration settings of the hypervisor **400**. In some implementations, the manager **412f** may be configured to authenticate the maintenance terminal and/or maintenance data and use the maintenance data only when the maintenance terminal and/or maintenance data has been authenticated successfully.

(28) In some implementations, only the manager **412f** may interact directly with the maintenance terminal, and any of the virtual machines **410a-e** and/or applications **412a-e** may interact with the maintenance terminal indirectly, through the manager **412f**. In such implementations, all maintenance data that is input into the primary SOM **310** may be received at the manager **412f**. In other words, among all components of the primary SOM **310** that are disposed between the virtual machine **410a** (and/or router **412a**) and the secondary SOM **320**, only the manager **412f** may receive maintenance data. Centralizing the transmission of maintenance data in this manner is advantageous because it may increase the security of the primary SOM **310** and/or GR **110**. The maintenance data, as noted above, may include an instruction to perform a software update or change a configuration setting of a specified one of the virtual machines **410a-f** and/or applications **412a-f**. The manager **412f** may communicate directly with the specified one of the virtual machines **410b-e** and/or applications **412b-e** to execute the instruction. In some implementations, the manager **412f** may use the secure shell (SSH) protocol to communicate with the specified one of the virtual machines **410b-e** and/or applications **412b-e** to execute the instruction. The “specified” one of the virtual machines **410a-f** and/or applications **412a-f** may be any of the virtual machines **410a-e** and/or applications **412a-e**. In some implementations, the manager **412f** may receive the maintenance data via a web interface that is provided by the manager **412f** and/or another component of the primary SOM **310**.

(29) The virtual machine **410c** may be configured to execute a VPN server **412c**. The VPN server **412c** may be configured to establish or maintain one or more VPN channels between GR **110** and devices in the wireless network **302**. More particularly, the VPN server **412c** may be configured to receive (from the router **312a**) packets that are designated for one or more devices in the wireless network **302**. Next, the VPN server **412c** may encrypt the received packets to produce encrypted packets. And finally, the VPN server **412c** may provide the encrypted packets to the virtual machine **410d** (and/or firewall **412d**).

(30) The virtual machine **410d** may be configured to execute a firewall application **412d** (hereinafter firewall **412d**). The firewall **412d** may be configured to monitor data traffic (e.g., one or more data packets) that is designated for one or more devices in the wireless network **302**. In operation, the firewall **412d** may receive encrypted packets from the VPN server **312c** and output the received packets (if they are not blocked) on data port **314**, thus causing the packets to be received at the secondary SOM **320**. The secondary SOM **320** may then transmit the encrypted packets over the wireless network **302** in a well-known fashion.

(31) Stated succinctly, data that is received at the GR **110** from the wired network **304** may travel across the primary SOM **310** along two data paths. Specifically, packets that are received from the wired network **304**, and which are designated for devices in the wireless network **302**, may pass through the router **412a**, the VPN server **412c**, and the firewall **412d**. On the other hand, packets that are designated for one of the virtual machines **410b-f** and/or applications **412b-f** may pass through the router **412a** and the virtual switch **420** before they reach their final destination. Under the nomenclature of the present disclosure, any device that is connected directly or indirectly to one of the data ports **312** of the GR **110** is considered to be part of the wired network **304**.

(32) In some implementations, when maintenance is performed on any of the virtual machines **410b-f** and/or applications **412b-f**, data that is generated by the virtual machines **410b-f** and/or applications **412b-f** (over the course of performing the maintenance) may be transmitted back to the

maintenance terminal that is overseeing the maintenance. Such data may pass through the virtual switch **420** and the router **412a** before it is transmitted over the wired network **304** (via one of the data ports **312**). In other words, the firewall **412b** may be configured to monitor data that originates from any of the virtual machines **410b-f** (and/or applications **412b-f**), over the course of performing maintenance, and forward the data to the router **412a**. As can be readily appreciated, in some implementations, the firewall **412b** may be dedicated to monitoring maintenance data that is transmitted to the manager **412f** and/or data that is generated, by one of the virtual machines **410a-f** (and/or applications **412a-f**) in response to the maintenance data. In such implementations, the firewall **412b** may not receive any user data that is being exchanged between networks **302** and **304** via the GR **110**. The router **412a** may then route the data to one of the data ports **312** for transmission over the wired network **304**.

(33) In some implementations, when the secondary SOM **320** receives data over the wireless network **302**, the secondary SOM **320** may provide the received data to the primary SOM **310**, and the primary SOM **310** may transmit the received data over the wired network **304**. In such situations, the data (received from the wireless network **302**) may pass through the firewall **412d**, the VPN server **412c**, and the router **412a** before it is forwarded to its final destination (in the wired network **304**) via one of the data ports **312**. More particularly, the firewall **412d** may monitor data that is received from the secondary SOM **320** (on data port **314**) and forward the received data (provided that it is not blocked) to the VPN server **412c**. The VPN server **412c** may decrypt data that is received from the firewall **412d** and forward the decrypted data to the router **412a**. And the router **412a** may route the decrypted data to one of the data ports **312** for transmission over the wired network **304**.

(34) The virtual machine **410e** may be configured to execute a wireless network authentication server **412e**. According to the present example, the server **412e** is a RADIUS server, and it may be arranged to provide centralized authentication, authorization, and management services for the wireless network **302**. As noted above, the server **412e** may be configured to authenticate devices that attempt to connect to the wireless network **302**. In this regard, when a device attempts to connect to the wireless network **302**, the device may provide an authentication credential (e.g., a PKI certificate) to the secondary SOM **320**. The secondary SOM **320** may transmit the authentication credential (via port **322**) to the primary SOM **310**. When the credential arrives at the primary SOM **310**, the credential may first stop at the firewall **412d** (via data port **314**). Afterwards, the firewall **412d** may forward the credential to the server **412e**. The server **412e** may attempt to authenticate the credential, after which the server **412e** may generate a response indicating whether the credential is valid. The firewall **412d** may receive the response from the server **412e** and forward the response to the secondary SOM **320**. If the response indicates that the credential is valid, the secondary SOM **320** may permit the device to connect to the wireless network **302**. Otherwise, if the response indicates that the credential is invalid, the secondary SOM **320** may deny the attempt to connect to the wireless network **302**.

(35) According to the example of FIG. 4, each of the virtual machines **410b-f** is executed in the volatile memory **318** of the primary SOM **310**, and its state is not persisted on the non-volatile memory **319** (or in any other non-volatile memory). In other words, under the present arrangement, each of the virtual machines is fully contained in the volatile memory **318**. More particularly, under the present arrangement, no state information of any of the virtual machines **410b-f** and/or applications **412b-d** is preserved after the GR is power-cycled (i.e., powered off and then powered on again). This arrangement is advantageous because it permits any intrusions into the primary SOM **310** to be interrupted by simply power-cycling the GR **110**. As can be readily appreciated, power-cycling the GR **110** would cause any malicious code (or other malicious data) to be destroyed when the GR **110** is power cycled. Moreover, power-cycling the GR **110** would cause the GR **110** to return to a default state in which the security of the GR **110** is fully intact.

(36) According to the example of FIG. 4, each of the virtual machines **410a-f** may be a lightweight

virtual machine having a memory footprint of less than 32 MB. In some respects, reducing the memory footprint of each of the virtual machines **410a-f** may allow the GR **110** to be power-cycled in less than 1 minute. In other words, the use of small-footprint virtual machines that are stored in volatile memory only allows the GR **110** to recover from cyber attacks in less than 1 minute. This is in contrast to some similar conventional systems, which may take days to recover.

(37) According to the example of FIG. **4**, none of the virtual machines **410b-f** may be configured (and/or permitted) to access the non-volatile memory **319**. Moreover, according to the example of FIG. **4**, only the virtual machine **410a** (out of virtual machines **410a-f**) may be configured (and/or permitted) to read and write data to the non-volatile memory **319**.

(38) According to the example of FIG. **4**, the virtual machines **410a-f** and/or applications **412a-f** may be connected to one another via virtual Ethernet links **416**, which are provided by the hypervisor **400**. In some respects, the Ethernet links **416a**, **416b**, and **416c** may form a first virtual network that is used for connecting the wireless network **302** and the wired network **304** to one another. Moreover, the link **416d** may form a second virtual network that is used for maintenance of the primary SOM **310**. In some implementations, the first virtual network may be used to carry data that crosses the primary SOM **310** en route to the wireless network **302** or wired network **304**. By contrast, the second virtual network may be used to carry maintenance data. In some respects, link **416c** may be configured to transmit unencrypted data; link **416b** may be configured to carry data that has been encrypted by the VPN server **412c**; link **416a** may be configured to carry authentication data (e.g., data that is formatted in accordance with the RADIUS protocol); and links **416d** may be configured to carry maintenance data.

(39) In some respects, the virtual machine **410a** (and/or applications **412a** and **414a**) may be part of a first security domain **422** of the primary SOM **310** and the virtual machines **410b-f** (and/or applications) may be part of a second security domain **424** of the primary SOM **310**. As discussed above, the second security domain **424** may be separated from the secondary SOM **320** (and/or wireless network **302**) by firewall **412d**. In addition, the second security domain **424** may also be separated from the first security domain **422** by firewall **412b**. In some respects, the use of firewalls to monitor all egress and ingress into the second security domain **424** can prevent any intrusions into one of the wireless network **302** (and/or secondary SOM **320**) and the first security domain **422** from spreading into the second security domain **424** of the primary SOM **310**. In some respects, the second security domain **424** may be operable to contain any malicious code (or other data) that enters the GR **110** through one of the wireless network **302** and the wired network **304**.

(40) FIG. **5** is a flowchart of an example of a process **500** that is performed by the primary SOM **310**, according to aspects of the disclosure. At step **502**, a boot loader of the primary SOM **310** starts the hypervisor **400**. At step **504**, the hypervisor **400** mounts a storage drive that is implemented in the non-volatile memory **319**. At step **506**, the hypervisor **400** instantiates a random-access memory (RAM) disk in the volatile memory **318**. At step **508**, the hypervisor **400** partitions the RAM disk into a plurality of disk partitions. Each of the partitions may include a different respective .ext2/4 file system and/or any other suitable type of file system instance. The RAM disk may include any suitable type of virtual drive that is instantiated in volatile memory.

(41) At step **510**, the hypervisor **400** identifies a plurality of virtual machine boot images that are stored in the storage drive that is mounted at step **504**). According to the present example, each of the virtual machine images may be an image of a different one of the virtual machines **410a-f** (and/or applications **412a-f**). At step **512**, the hypervisor **400** boots each of the virtual machine images on a different one of the RAM disk partitions (obtained at step **508**). As a result of executing step **512**, each of the virtual machines **410a-e** begins executing on a different RAM disk partition. At step **514**, the hypervisor **400** gives the virtual machine **410a** control over the non-volatile memory **319** or the storage device that is implemented in the non-volatile memory **319**. Giving control over the non-volatile memory **319** may include changing one or more configuration settings of the hypervisor to allow the virtual machine **410a** to read and write data to the non-

volatile memory **319** (or storage drive that is implemented on the non-volatile memory **319**). Additionally or alternatively, giving control over the non-volatile memory **319** may include providing the virtual machine **410a** with an address (e.g., a PCI address) of the non-volatile memory **319** and/or any other information that is needed (by the virtual machine **410a**) to access the non-volatile memory **319**.

(42) FIG. **6** is a state diagram illustrating aspects of the operation of the GR **110**. As illustrated, throughout its operation, the GR **110** may be in one of an inactive state **610**, a boot state **620**, a start-up state **630**, and an active state **640**.

(43) When the GR **110** is in the boot state, the GR **110** may be powered off or otherwise inoperative. The GR **110** may exit the inactive state **610** and transition to the boot state **620** when the GR **110** is powered on or otherwise turned on.

(44) Upon entering the boot state **620**, a bootloader of the GR **110** may start the hypervisor **400**. Next, the hypervisor **400** may instantiate a RAM disk in the volatile memory **318**. After the RAM disk is instantiated, the GR **110** may transition into the start-up state **630**.

(45) Upon entering the start-up state **630**, the hypervisor **400** may partition the RAM disk into a plurality of partitions. As noted above, each partition may have a separate file system. Next, the hypervisor **400** may boot each of the virtual machines **410a-f** on a different one of the partitions. After each of the virtual machines **410** is booted, the GR **110** may transition into the active state **640**.

(46) When the GR **110** is in the active state **640**, the GR **110** may route data traffic from the wireless network **302** to the wired network **304**, and vice versa. The GR **110** may transition out of the operating state **640** and back into the inactive state **610** when the GR **110** is powered off. When the GR **110** is powered off, all data that is stored in the RAM disk partitions will be destroyed, causing the reversal of any intrusions into one or more of the virtual machines **410a-e**.

(47) FIG. **7** shows a perspective side view of the GR **110** according to aspects of the disclosure. As illustrated, the GR **110** may include a housing **701**, a plurality of connectors **702**, a plurality of connectors **704**, a power switch **706**, and a purge switch **708**. As noted above, each of the connectors **702** may be arranged to connect the GR **110** to an antenna for transmitting and receiving data over the wireless network **302**. Each of the connectors **704** may be arranged to connect to one or more Ethernet lines. The power switch **706** may be arranged to power the GR **110** on and off, as well as restart the GR **110**. The purge switch **708** may be configured to erase all contents of the volatile memory **318** when activated. In some implementations, upon detecting the purge switch **708** is activated, the GR **110** may cut all supply of power to the volatile memory **318**, thereby causing the partitions used to execute the virtual machines **410** to be destroyed. As noted above, destroying the partitions may cause all code that is used to execute the virtual machines to be deleted. Additionally or alternatively, in some implementations, upon detecting the purge switch **708** is activated, the GR **110** may execute a series of overwrite sequences on all volatile and non-volatile memories to render them unrecoverable. In other words, executing the series of overwrite sequences may prevent data of any of the virtual machines **410a-f** (and/or applications **412a-f**) from being recovered from the volatile memory **318**.

(48) In some implementations, the housing **701** may have the dimensions of 9.5 in×5.5 in×5 in. In this respect, FIG. **7** is provided to illustrate that the GR **110** (in some implementations) can be a small form factor, rugged, tactical computer that integrates all of the security applications and server capabilities that are typically met by an entire data center. Instead of racks of servers and infrastructure components that are typically required to serve hundreds or thousands of users at the enterprise level, the GR **110** may provide all of these capabilities to tens of users in a ground vehicle or aviation environment. In some respects, the switch **330** may be a IO-port gigabit Ethernet switch. Furthermore, the GR **110** may operate as a 5-port router, a dual authentication server, a secure wireless access point, a firewall, an intrusion detection system, and a software appliance (e.g., see user application **114a**)—while having a relatively small footprint in comparison

to similar systems. In some respects, the small form-factor of the GR **110** is achieved by using virtualization to implement different components of the first security domain **422** and the second security domain **424** of the primary SOM **310**. Although the GR **110** is described as suitable for use in mobile environments, it will be understood that the GR **110** may also be used in more traditional settings where WiFi routers are used, such as wireless networks of university campuses and office buildings for example.

(49) In some implementations, each of the primary SOM **310** and the secondary SOM **320** may be implemented as a separate integrated circuit (e.g., chip), and they may be mounted on the same circuit board. Additionally or alternatively, in some implementations, the primary SOM **310** and the secondary SOM **320** may be implemented as separate circuit boards that are coupled to one another. Stated succinctly, the present disclosure is not limited to any specific implementation of the primary SOM **310** and the secondary SOM **320**.

(50) Although in the example of FIG. **4** the virtual machines **410a-f** and/or applications **412a-f** are connected to one another via Ethernet data links, it will be understood that the present disclosure is not limited to any specific method for connecting the virtual machines **410a-f** and/or applications **412a-f** to one another. Throughout the disclosure the term “application” is used to refer to any of the router **412a**, the firewall **412b**, the VPN server **412c**, the firewall **412d**, the wireless network authentication server **412e**, and the manager **412e**. However, it will be understood that the use of this term is in no way intended to imply a specific architecture for the virtual machines **412a-f** and/or applications **412a-f**. For example, in some implementations, any of the router **412a**, the firewall **412b**, the VPN server **412c**, the firewall **412d**, the wireless network authentication server **412e**, and the manager **412e** may be integrated into its respective virtual machine (e.g., as part of a kernel of the virtual machine). As another example, any of the router **412a**, the firewall **412b**, the VPN server **412c**, the firewall **412d**, the wireless network authentication server **412e**, and the manager **412e** may be provided separately of its virtual machine (e.g., as part of a service or user application that is executed on top of a kernel of the virtual machine). Stated succinctly, the present disclosure is not limited to any specific implementation of the virtual machines **410a-f** and/or applications **412a-f**.

(51) According to the example of FIG. **3**, the wireless network **302** is the same or similar to any of the wireless networks **120** and **220**, which are shown in FIGS. **1** and **2**, respectively. In this regard, it will be understood that the present disclosure is not limited to any specific implementation of the wireless network **302**. According to the example of FIG. **3**, the wired network **304** is the same or similar to any of the wired networks **130** and **230**, which are shown in FIGS. **1** and **2**, respectively. In this regard, it will be understood that the present disclosure is not limited to any specific implementation of the wired network **304**. According to the example of FIG. **4**, the hypervisor **400** is a bare metal hypervisor. However, alternative implementations are possible in which the hypervisor **400** is a hosted hypervisor. Stated succinctly, it will be understood that the present disclosure is not limited to any specific implementations of the hypervisor.

(52) According to the example of FIG. **3** each of the virtual machines **410a-b** includes a custom-built Linux distribution image that includes as few services as possible. Each of the applications **412a-b** is executed as part of (or on top) of the Linux kernel of its respective virtual machine. However, as noted above, the present disclosure is not limited to any specific implementation of the virtual machines **410a-f** and/or applications **412a-f**.

(53) The processes described herein may be implemented in hardware, software, or a combination of the two. The processes described herein may be implemented in computer programs executed on programmable computers/machines that each includes a processor, a non-transitory machine-readable medium or another article of manufacture that is readable by the processor (including volatile and non-volatile memory and/or storage elements), at least one input device, and one or more output devices. Program code may be applied to data entered using an input device to perform any of the processes described herein and to generate output information.

(54) The system may be implemented, at least in part, via a computer program product, (e.g., in a non-transitory machine-readable storage medium such as, for example, a non-transitory computer-readable medium), for execution by, or to control the operation of, data processing apparatus (e.g., a programmable processor, a computer, or multiple computers). Each such program may be implemented in a high-level procedural or object-oriented programming language to work with the rest of the computer-based system. However, the programs may be implemented in assembly, machine language, or Hardware Description Language. The language may be a compiled or an interpreted language, and it may be deployed in any form, including as a stand-alone program or as a module, component, subroutine, or another unit suitable for use in a computing environment. A computer program may be deployed to be executed on one computer or multiple computers at one site or distributed across multiple sites and interconnected by a communication network. A computer program may be stored on a non-transitory machine-readable medium that is readable by a general or special purpose programmable computer for configuring and operating the computer when the non-transitory machine-readable medium is read by the computer to perform the processes described herein. For example, the processes described herein may also be implemented as a non-transitory machine-readable storage medium, configured with a computer program, where upon execution, instructions in the computer program cause the computer to operate in accordance with the processes. A non-transitory machine-readable medium may include but is not limited to a hard drive, compact disc, flash memory, non-volatile memory, volatile memory, magnetic diskette and so forth but does not include a transitory signal per se.

(55) Having described preferred embodiments, which serve to illustrate various concepts, structures and techniques, which are the subject of this patent, it will now become apparent that other embodiments incorporating these concepts, structures and techniques may be used. Accordingly, it is submitted that the scope of the patent should not be limited to the described embodiments but rather should be limited only by the spirit and scope of the following claims.

Claims

1. An apparatus comprising: a non-volatile memory; a volatile memory; a first electronic circuit that is configured to operate as a wireless access point, the first electronic circuit including a wireless controller for accessing a wireless network; and a second electronic circuit that is operatively coupled to the first electronic circuit, the second electronic circuit including at least one processor configured to execute: (i) a first virtual machine that includes a virtual private network (VPN) server, and (ii) a second virtual machine that includes a first firewall, and wherein the VPN server is configured to (i) encrypt data that is received at the apparatus to produce encrypted data, and (ii) forward the encrypted data to the wireless controller for transmission over the wireless network.
2. The apparatus of claim 1, wherein each of the first virtual machine and the second virtual machine is fully contained in the volatile memory.
3. The apparatus of claim 1, wherein the second electronic circuit is operatively coupled to the first electronic circuit via a wired connection.
4. The apparatus of claim 1, wherein the second electronic circuit is operatively coupled to the first electronic circuit via a wireless connection.
5. The apparatus of claim 1, wherein the first firewall is interposed between the wireless controller and the VPN server, and arranged to monitor traffic between the VPN server and the wireless controller, the first electronic circuit includes a first system on module (SOM), and the second electronic circuit includes a second SOM.
6. The apparatus of claim 1, wherein the processor is configured to execute a third virtual machine, the third virtual machine being configured to execute a wireless network authentication server that is operatively coupled to the wireless controller via the first firewall, the wireless network

authentication server being configured to authenticate devices that attempt to join the wireless network.

7. The apparatus of claim 1, wherein: the first electronic circuit and the second electronic circuit are coupled to one another via an Ethernet connection, and the encrypted data is forwarded to the second electronic circuit via the Ethernet connection; and the processor is further configured to execute a hypervisor, the hypervisor being arranged to perform the operations of: instantiating a random-access memory (RAM) disk in the volatile memory, partitioning the RAM disk into a plurality of partitions, and launching each of the first virtual machine and the second virtual machine on a different one of the plurality of partitions.

8. The apparatus of claim 1, wherein the processor is configured to execute a third virtual machine, the third virtual machine being configured to execute a manager application, the manager application being arranged to change a configuration of at least one of the VPN server and the first firewall based on maintenance data that is received at the apparatus.

9. The apparatus of claim 8, wherein the processor is further configured to execute a fourth virtual machine, the fourth virtual machine being configured to execute a router that is arranged to route data that is received at the apparatus to one of the VPN server, external Ethernet ports, and the manager application.

10. The apparatus of claim 1, wherein the first virtual machine is executed in a first partition that is instantiated in the volatile memory, and the second virtual machine is executed in a second partition that is instantiated in the volatile memory, each of the first partition and the second partition having a separate file system.

11. A method for use in an electronic device that includes a first electronic circuit configured to operate as an access point for accessing a wireless network, and a second electronic circuit having a volatile memory and at least one processor, the method comprising: instantiating a random-access memory (RAM) disk in the volatile memory of the second electronic circuit; partitioning the RAM disk into a plurality of partitions; launching a first virtual machine on the second electronic circuit, the first virtual machine (i) being launched in a first partition of the RAM disk, and (ii) including a virtual private network (VPN) server; and launching a second virtual machine on the second electronic circuit, the second virtual machine being launched in a second partition of the RAM disk, the second virtual machine including a first firewall.

12. The method of claim 11, wherein the VPN server is configured to (i) encrypt data that is received at the apparatus to produce encrypted data, and (ii) forward the encrypted data to the wireless controller for transmission over the wireless network.

13. The method of claim 11, further comprising: detecting that a purge switch of the electronic device is activated; and executing one or more overwrite sequences on the volatile memory of the second electronic circuit to render data of the VPN server unrecoverable.

14. The method of claim 11, further comprising: launching a hypervisor on the second electronic circuit when the electronic device is booted, wherein the RAM disk is instantiated and partitioned by the hypervisor, and wherein the first virtual machine and the second virtual machine are launched by the hypervisor.

15. The method of claim 11, wherein the first firewall is interposed between the wireless controller and the VPN server and arranged to monitor traffic between the VPN server and the wireless controller.

16. The method of claim 11, further comprising: launching a third virtual machine on the second electronic circuit, the third virtual machine being configured to execute a manager application, the manager application being arranged to change a configuration of at least one of the VPN server and the first firewall based on maintenance data that is received at the apparatus.

17. The method of claim 16, further comprising: launching a fourth virtual machine on the second electronic circuit, the fourth virtual machine being launched in a fourth partition of the RAM disk, the fourth virtual machine including a router configured to route data that is received at the

apparatus to one of the manager application, external Ethernet ports, and the VPN server.

18. The method of claim 11, wherein each of the first virtual machine and the second virtual machine is fully contained in the volatile memory.

19. The method of claim 11, wherein the second electronic circuit is operatively coupled to the first electronic circuit via a wired connection.

20. The method of claim 11, wherein the second electronic circuit is operatively coupled to the first electronic circuit via a wireless connection.
