

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12418513
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Du; Fan et al.

Internet protocol based security over port forwarding tunnels

Abstract

An application server host may receive, via a port forwarding tunnel, a connection request that includes an indication of a client internet protocol (IP) address for a client host. The application server host may identify based at least in part on the client IP address and using a mapping maintained for a plurality of client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The application server host may open a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

Inventors: Du; Fan (Santa Clara, CA), Luo; Jiangbin (Cupertino, CA), Bansal; Pradeep (Portland, OR), Jang; Keon (Los Altos, CA)

Applicant: Rubrik, Inc. (Palo Alto, CA)

Family ID: 1000008821423

Assignee: Rubrik, Inc. (Palo Alto, CA)

Appl. No.: 17/855764

Filed: June 30, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240007441 A1	Jan. 04, 2024

Publication Classification

Int. Cl.: H04L9/40 (20220101)

U.S. Cl.:

CPC **H04L63/029** (20130101); **H04L63/0236** (20130101); **H04L63/20** (20130101);

Field of Classification Search

CPC: H04L (63/029); H04L (63/0236); H04L (63/20); H04L (9/40)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9143480	12/2014	Brousseau	N/A	H04L 67/141
9172559	12/2014	Chen	N/A	H04L 12/4641
10375025	12/2018	Zheng	N/A	H04L 63/166
2002/0138628	12/2001	Tingley	709/227	H04L 61/00
2014/0181248	12/2013	Deutsch	709/217	H04L 61/5076
2017/0223154	12/2016	Hammam	N/A	H04L 45/60
2018/0124198	12/2017	Petrov	N/A	H04L 12/4633
2020/0213151	12/2019	Srivatsan	N/A	H04L 9/0643
2023/0283608	12/2022	Bosch	726/4	H04L 63/102

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2004007671	12/2003	JP	H04L 12/4633
WO-02061599	12/2001	WO	H04L 29/12009
WO-2009055722	12/2008	WO	H04L 29/12367

Primary Examiner: Jhaveri; Jayesh M

Attorney, Agent or Firm: Holland & Hart LLP

Background/Summary

FIELD OF TECHNOLOGY

(1) The present disclosure relates generally to database systems and data processing, and more specifically to internet protocol based security over port forwarding tunnels.

BACKGROUND

(2) A data management system (DMS) may be employed to manage data associated with one or more computing systems. The data may be generated, stored, or otherwise used by the one or more computing systems, examples of which may include servers, databases, virtual machines, cloud computing systems, file systems (e.g., network-attached storage (NAS) systems), or other data storage or processing systems. The DMS may provide data backup, data recovery, data classification, or other types of data management services for data of the one or more computing

systems. Improved data management may offer improved performance with respect to reliability, speed, efficiency, scalability, security, or ease-of-use, among other possible aspects of performance.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 illustrates an example of a data management system that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (2) FIG. 2 illustrates an example of a computing environment that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (3) FIG. 3 illustrates an example of a computing environment for multiple client networks that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (4) FIG. 4 illustrates an example of a process flow that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (5) FIG. 5 shows a block diagram of an apparatus that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (6) FIG. 6 shows a block diagram of a connection manager that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (7) FIG. 7 shows a diagram of a system including a device that supports internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.
- (8) FIGS. 8 through 12 show flowcharts illustrating methods that support internet protocol based security over port forwarding tunnels in accordance with aspects of the present disclosure.

DETAILED DESCRIPTION

- (9) Access to various computing resources, such as an application server, may be protected using security measures. For example, some application servers enforce security policies based on client internet protocol (IP) addresses, as the client IP address is generally coupled to a client identity. However, some services provide secure access to remote clients using port forwarding tunnels, which may replace the client IP address (of a packet) with the IP address of the port forwarding tunnel. In such cases, IP addresses of the client may not be used to provide secure access. Additionally, as some services may support multitenancy and client IP addresses across tenants may be the same, it may be difficult to differentiate between clients to enforce security policies.
- (10) Implementations described herein support secure access using client IP addresses and port forwarding tunnel connection techniques. When a new client is added to the system, the client is assigned a unique virtual IP address that is maintained on the application server host in a mapping (e.g., file or database). The unique virtual IP address is associated with a client network name and a client host IP address in the mapping. When a connection request is received at the application server host, the application server host identifies a corresponding virtual IP address and opens the connection using the identified virtual IP address. The virtual IP address may then be used by the application server to enforce security policies.
- (11) Aspects of the disclosure are initially described in the context of an environment supporting data management. Aspects of the disclosure are further described with respect to computing environments illustrating resource connections using port forwarding tunnels and a process flow diagram illustrating IP based security over port forwarding tunnels. Aspects of the disclosure are further illustrated by and described with reference to apparatus diagrams, system diagrams, and flowcharts that relate to IP based security over port forwarding tunnels.
- (12) FIG. 1 illustrates an example of a computing environment **100** supports IP based security over port forwarding tunnels in accordance with various aspects of the present disclosure. The computing environment **100** may include a computing system **105**, a data management system

(DMS) **110**, and one or more computing devices **115**, which may be in communication with one another via a network **120**. The computing system **105** may generate, store, process, modify, or otherwise use associated data, and the DMS **110** may provide one or more data management services for the computing system **105**. For example, the DMS **110** may provide a data backup service, a data recovery service, a data classification service, a data transfer or replication service, one or more other data management services, or any combination thereof for data associated with the computing system **105**.

(13) The network **120** may allow the one or more computing devices **115**, the computing system **105**, and the DMS **110** to communicate (e.g., exchange information) with one another. The network **120** may include aspects of one or more wired networks (e.g., the Internet), one or more wireless networks (e.g., cellular networks), or any combination thereof. The network **120** may include aspects of one or more public networks or private networks, as well as secured or unsecured networks, or any combination thereof. The network **120** also may include any quantity of communications links and any quantity of hubs, bridges, routers, switches, ports or other physical or logical network components.

(14) A computing device **115** may be used to input information to or receive information from the computing system **105**, the DMS **110**, or both. For example, a user of the computing device **115** may provide user inputs via the computing device **115**, which may result in commands, data, or any combination thereof being communicated via the network **120** to the computing system **105**, the DMS **110**, or both. Additionally, or alternatively, a computing device **115** may output (e.g., display) data or other information received from the computing system **105**, the DMS **110**, or both. A user of a computing device **115** may, for example, use the computing device **115** to interact with one or more user interfaces (e.g., graphical user interfaces (GUIs)) to operate or otherwise interact with the computing system **105**, the DMS **110**, or both. Though one computing device **115** is shown in FIG. 1, it is to be understood that the computing environment **100** may include any quantity of computing devices **115**.

(15) A computing device **115** may be a stationary device (e.g., a desktop computer or access point) or a mobile device (e.g., a laptop computer, tablet computer, or cellular phone). In some examples, a computing device **115** may be a commercial computing device, such as a server or collection of servers. And in some examples, a computing device **115** may be a virtual device (e.g., a virtual machine). Though shown as a separate device in the example computing environment of FIG. 1, it is to be understood that in some cases a computing device **115** may be included in (e.g., may be a component of) the computing system **105** or the DMS **110**.

(16) The computing system **105** may include one or more servers **125** and may provide (e.g., to the one or more computing devices **115**) local or remote access to applications, databases, or files stored within the computing system **105**. The computing system **105** may further include one or more data storage devices **130**. Though one server **125** and one data storage device **130** are shown in FIG. 1, it is to be understood that the computing system **105** may include any quantity of servers **125** and any quantity of data storage devices **130**, which may be in communication with one another and collectively perform one or more functions ascribed herein to the server **125** and data storage device **130**.

(17) A data storage device **130** may include one or more hardware storage devices operable to store data, such as one or more hard disk drives (HDDs), magnetic tape drives, solid-state drives (SSDs), storage area network (SAN) storage devices, or network-attached storage (NAS) devices. In some cases, a data storage device **130** may comprise a tiered data storage infrastructure (or a portion of a tiered data storage infrastructure). A tiered data storage infrastructure may allow for the movement of data across different tiers of the data storage infrastructure between higher-cost, higher-performance storage devices (e.g., SSDs and HDDs) and relatively lower-cost, lower-performance storage devices (e.g., magnetic tape drives). In some examples, a data storage device **130** may be a database (e.g., a relational database), and a server **125** may host (e.g., provide a database

management system for) the database.

(18) A server **125** may allow a client (e.g., a computing device **115**) to download information or files (e.g., executable, text, application, audio, image, or video files) from the computing system **105**, to upload such information or files to the computing system **105**, or to perform a search query related to particular information stored by the computing system **105**. In some examples, a server **125** may act as an application server or a file server. In general, a server **125** may refer to one or more hardware devices that act as the host in a client-server relationship or a software process that shares a resource with or performs work for one or more clients.

(19) A server **125** may include a network interface **140**, processor **145**, memory **150**, disk **155**, and computing system manager **160**. The network interface **140** may enable the server **125** to connect to and exchange information via the network **120** (e.g., using one or more network protocols). The network interface **140** may include one or more wireless network interfaces, one or more wired network interfaces, or any combination thereof. The processor **145** may execute computer-readable instructions stored in the memory **150** in order to cause the server **125** to perform functions ascribed herein to the server **125**. The processor **145** may include one or more processing units, such as one or more central processing units (CPUs), one or more graphics processing units (GPUs), or any combination thereof. The memory **150** may comprise one or more types of memory (e.g., random access memory (RAM), static random access memory (SRAM), dynamic random access memory (DRAM), read-only memory ((ROM), electrically erasable programmable read-only memory (EEPROM), Flash, etc.). Disk **155** may include one or more HDDs, one or more SSDs, or any combination thereof. Memory **150** and disk **155** may comprise hardware storage devices. The computing system manager **160** may manage the computing system **105** or aspects thereof (e.g., based on instructions stored in the memory **150** and executed by the processor **145**) to perform functions ascribed herein to the computing system **105**. In some examples, the network interface **140**, processor **145**, memory **150**, and disk **155** may be included in a hardware layer of a server **125**, and the computing system manager **160** may be included in a software layer of the server **125**. In some cases, the computing system manager **160** may be distributed across (e.g., implemented by) multiple servers **125** within the computing system **105**.

(20) In some examples, the computing system **105** or aspects thereof may be implemented within one or more cloud computing environments, which may alternatively be referred to as cloud environments. Cloud computing may refer to Internet-based computing, wherein shared resources, software, and/or information may be provided to one or more computing devices on-demand via the Internet. A cloud environment may be provided by a cloud platform, where the cloud platform may include physical hardware components (e.g., servers) and software components (e.g., operating system) that implement the cloud environment. A cloud environment may implement the computing system **105** or aspects thereof through Software-as-a-Service (SaaS) or Infrastructure-as-a-Service (IaaS) services provided by the cloud environment. SaaS may refer to a software distribution model in which applications are hosted by a service provider and made available to one or more client devices over a network (e.g., to one or more computing devices **115** over the network **120**). IaaS may refer to a service in which physical computing resources are used to instantiate one or more virtual machines, the resources of which are made available to one or more client devices over a network (e.g., to one or more computing devices **115** over the network **120**).

(21) In some examples, the computing system **105** or aspects thereof may implement or be implemented by one or more virtual machines. The one or more virtual machines may run various applications, such as a database server, an application server, or a web server. For example, a server **125** may be used to host (e.g., create, manage) one or more virtual machines, and the computing system manager **160** may manage a virtualized infrastructure within the computing system **105** and perform management operations associated with the virtualized infrastructure. The computing system manager **160** may manage the provisioning of virtual machines running within the virtualized infrastructure and provide an interface to a computing device **115** interacting with the

virtualized infrastructure. For example, the computing system manager **160** may be or include a hypervisor and may perform various virtual machine-related tasks, such as cloning virtual machines, creating new virtual machines, monitoring the state of virtual machines, moving virtual machines between physical hosts for load balancing purposes, and facilitating backups of virtual machines. In some examples, the virtual machines, the hypervisor, or both, may virtualize and make available resources of the disk **155**, the memory, the processor **145**, the network interface **140**, the data storage device **130**, or any combination thereof in support of running the various applications. Storage resources (e.g., the disk **155**, the memory **150**, or the data storage device **130**) that are virtualized may be accessed by applications as a virtual disk.

(22) The DMS **110** may provide one or more data management services for data associated with the computing system **105** and may include DMS manager **190** and any quantity of storage nodes **185**. The DMS manager **190** may manage operation of the DMS **110**, including the storage nodes **185**. Though illustrated as a separate entity within the DMS **110**, the DMS manager **190** may in some cases be implemented (e.g., as a software application) by one or more of the storage nodes **185**. In some examples, the storage nodes **185** may be included in a hardware layer of the DMS **110**, and the DMS manager **190** may be included in a software layer of the DMS **110**. In the example illustrated in FIG. 1, the DMS **110** is separate from the computing system **105** but in communication with the computing system **105** via the network **120**. It is to be understood, however, that in some examples at least some aspects of the DMS **110** may be located within computing system **105**. For example, one or more servers **125**, one or more data storage devices **130**, and at least some aspects of the DMS **110** may be implemented within the same cloud environment or within the same data center.

(23) Storage nodes **185** of the DMS **110** may include respective network interfaces **165**, processors **170**, memories **175**, and disks **180**. The network interfaces **165** may enable the storage nodes **185** to connect to one another, to the network **120**, or both. A network interface **165** may include one or more wireless network interfaces, one or more wired network interfaces, or any combination thereof. The processor **170** of a storage node **185** may execute computer-readable instructions stored in the memory **175** of the storage node **185** in order to cause the storage node **185** to perform processes described herein as performed by the storage node **185**. A processor **170** may include one or more processing units, such as one or more CPUs, one or more GPUs, or any combination thereof. The memory **150** may comprise one or more types of memory (e.g., RAM, SRAM, DRAM, ROM, EEPROM, Flash, etc.). A disk **180** may include one or more HDDs, one or more SSDs, or any combination thereof. Memories **175** and disks **180** may comprise hardware storage devices. Collectively, the storage nodes **185** may in some cases be referred to as a storage cluster or as a cluster of storage nodes **185**.

(24) The DMS **110** may provide a backup and recovery service for the computing system **105**. For example, the DMS **110** may manage the extraction and storage of snapshots **135** associated with different point-in-time versions of one or more target computing objects within the computing system **105**. A snapshot **135** of a computing object (e.g., a virtual machine, a database, a filesystem, a virtual disk, a virtual desktop, or other type of computing system or storage system) may be a file (or set of files) that represents a state of the computing object (e.g., the data thereof) as of a particular point in time. A snapshot **135** may also be used to restore (e.g., recover) the corresponding computing object as of the particular point in time corresponding to the snapshot **135**. A computing object of which a snapshot **135** may be generated may be referred to as snappable. Snapshots **135** may be generated at different times (e.g., periodically or on some other scheduled or configured basis) in order to represent the state of the computing system **105** or aspects thereof as of those different times. In some examples, a snapshot **135** may include metadata that defines a state of the computing object as of a particular point in time. For example, a snapshot **135** may include metadata associated with (e.g., that defines a state of) some or all data blocks included in (e.g., stored by or otherwise included in) the computing object. Snapshots **135** (e.g.,

collectively) may capture changes in the data blocks over time. Snapshots **135** generated for the target computing objects within the computing system **105** may be stored in one or more storage locations (e.g., the disk **155**, memory **150**, the data storage device **130**) of the computing system **105**, in the alternative or in addition to being stored within the DMS **110**, as described below.

(25) To obtain a snapshot **135** of a target computing object associated with the computing system **105** (e.g., of the entirety of the computing system **105** or some portion thereof, such as one or more databases, virtual machines, or filesystems within the computing system **105**), the DMS manager **190** may transmit a snapshot request to the computing system manager **160**. In response to the snapshot request, the computing system manager **160** may set the target computing object into a frozen state (e.g., a read-only state). Setting the target computing object into a frozen state may allow a point-in-time snapshot **135** of the target computing object to be stored or transferred.

(26) In some examples, the computing system **105** may generate the snapshot **135** based on the frozen state of the computing object. For example, the computing system **105** may execute an agent of the DMS **110** (e.g., the agent may be software installed at and executed by one or more servers **125**), and the agent may cause the computing system **105** to generate the snapshot **135** and transfer the snapshot to the DMS **110** in response to the request from the DMS **110**. In some examples, the computing system manager **160** may cause the computing system **105** to transfer, to the DMS **110**, data that represents the frozen state of the target computing object, and the DMS **110** may generate a snapshot **135** of the target computing object based on the corresponding data received from the computing system **105**.

(27) Once the DMS **110** receives, generates, or otherwise obtains a snapshot **135**, the DMS **110** may store the snapshot **135** at one or more of the storage nodes **185**. The DMS **110** may store a snapshot **135** at multiple storage nodes **185**, for example, for improved reliability. Additionally, or alternatively, snapshots **135** may be stored in some other location connected with the network **120**. For example, the DMS **110** may store more recent snapshots **135** at the storage nodes **185**, and the DMS **110** may transfer less recent snapshots **135** via the network **120** to a cloud environment (which may include or be separate from the computing system **105**) for storage at the cloud environment, a magnetic tape storage device, or another storage system separate from the DMS **110**.

(28) Updates made to a target computing object that has been set into a frozen state may be written by the computing system **105** to a separate file (e.g., an update file) or other entity within the computing system **105** while the target computing object is in the frozen state. After the snapshot **135** (or associated data) of the target computing object has been transferred to the DMS **110**, the computing system manager **160** may release the target computing object from the frozen state, and any corresponding updates written to the separate file or other entity may be merged into the target computing object.

(29) In response to a restore command (e.g., from a computing device **115** or the computing system **105**), the DMS **110** may restore a target version (e.g., corresponding to a particular point in time) of a computing object based on a corresponding snapshot **135** of the computing object. In some examples, the corresponding snapshot **135** may be used to restore the target version based on data of the computing object as stored at the computing system **105** (e.g., based on information included in the corresponding snapshot **135** and other information stored at the computing system **105**, the computing object may be restored to its state as of the particular point in time). Additionally, or alternatively, the corresponding snapshot **135** may be used to restore the data of the target version based on data of the computing object as included in one or more backup copies of the computing object (e.g., file-level backup copies or image-level backup copies). Such backup copies of the computing object may be generated in conjunction with or according to a separate schedule than the snapshots **135**. For example, the target version of the computing object may be restored based on the information in a snapshot **135** and based on information included in a backup copy of the target object generated prior to the time corresponding to the target version. Backup copies of the

computing object may be stored at the DMS **110** (e.g., in the storage nodes **185**) or in some other location connected with the network **120** (e.g., in a cloud environment, which in some cases may be separate from the computing system **105**).

(30) In some examples, the DMS **110** may restore the target version of the computing object and transfer the data of the restored computing object to the computing system **105**. And in some examples, the DMS **110** may transfer one or more snapshots **135** to the computing system **105**, and restoration of the target version of the computing object may occur at the computing system **105** (e.g., as managed by an agent of the DMS **110**, where the agent may be installed and operate at the computing system **105**).

(31) In response to a mount command (e.g., from a computing device **115** or the computing system **105**), the DMS **110** may instantiate data associated with a point-in-time version of a computing object based on a snapshot **135** corresponding to the computing object (e.g., along with data included in a backup copy of the computing object) and the point-in-time. The DMS **110** may then allow the computing system **105** to read or modify the instantiated data (e.g., without transferring the instantiated data to the computing system). In some examples, the DMS **110** may instantiate (e.g., virtually mount) some or all of the data associated with the point-in-time version of the computing object for access by the computing system **105**, the DMS **110**, or the computing device **115**.

(32) In some examples, the DMS **110** may store different types of snapshots, including for the same computing object. For example, the DMS **110** may store both base snapshots **135** and incremental snapshots **135**. A base snapshot **135** may represent the entirety of the state of the corresponding computing object as of a point in time corresponding to the base snapshot **135**. An incremental snapshot **135** may represent the changes to the state—which may be referred to as the delta—of the corresponding computing object that have occurred between an earlier or later point in time corresponding to another snapshot **135** (e.g., another base snapshot **135** or incremental snapshot **135**) of the computing object and the incremental snapshot **135**. In some cases, some incremental snapshots **135** may be forward-incremental snapshots **135** and other incremental snapshots **135** may be reverse-incremental snapshots **135**. To generate a full snapshot **135** of a computing object using a forward-incremental snapshot **135**, the information of the forward-incremental snapshot **135** may be combined with (e.g., applied to) the information of an earlier base snapshot **135** of the computing object along with the information of any intervening forward-incremental snapshots **135**, where the earlier base snapshot **135** may include a base snapshot **135** and one or more reverse-incremental or forward-incremental snapshots **135**. To generate a full snapshot **135** of a computing object using a reverse-incremental snapshot **135**, the information of the reverse-incremental snapshot **135** may be combined with (e.g., applied to) the information of a later base snapshot **135** of the computing object along with the information of any intervening reverse-incremental snapshots **135**.

(33) In some examples, the DMS **110** may provide a data classification service, a malware detection service, a data transfer or replication service, backup verification service, or any combination thereof, among other possible data management services for data associated with the computing system **105**. For example, the DMS **110** may analyze data included in one or more computing objects of the computing system **105**, metadata for one or more computing objects of the computing system **105**, or any combination thereof, and based on such analysis, the DMS **110** may identify locations within the computing system **105** that include data of one or more target data types (e.g., sensitive data, such as data subject to privacy regulations or otherwise of particular interest) and output related information (e.g., for display to a user via a computing device **115**). Additionally, or alternatively, the DMS **110** may detect whether aspects of the computing system **105** have been impacted by malware (e.g., ransomware). Additionally, or alternatively, the DMS **110** may relocate data or create copies of data based on using one or more snapshots **135** to restore the associated computing object within its original location or at a new location (e.g., a new

location within a different computing system **105**). Additionally, or alternatively, the DMS **110** may analyze backup data to ensure that the underlying data (e.g., user data or metadata) has not been corrupted. The DMS **110** may perform such data classification, malware detection, data transfer or replication, or backup verification, for example, based on data included in snapshots **135** or backup copies of the computing system **105**, rather than live contents of the computing system **105**, which may beneficially avoid adversely.

(34) As described herein, the server **125** may run or host various applications (e.g., via one or more virtual machines), such as an application server, a database server or, a webserver. In some examples, a client system, such as computing device **115** or the DMS **110**, may access resources of the computing system **105** and the server **125** or may store client data to the resources of the computing system **105**. Similarly, the DMS **110** may function as a server that is accessible by client systems, such as the computing device **115** or the computing system **105**. As such, computing system **105**, DMS **110**, or another system (e.g., a cloud computing resource) may function as a server that supports access to resources and data by client hosts.

(35) In some examples, a client system may access resources of a server over tunnels. In one example, a port forwarding tunnel technique, such as secure shell protocol (SSH) and secure socket funneling (SSF) may be used to provide secure access to remote clients to resources of a server. Via a tunnel, the client systems may access remote resources and services, such as file access, database access, among others. Many services, such as network file system (NFS) and server message block (SMB) may enforce security policies based on a source IP address of the client. Thus, if the client IP address is not authorized, then access to services and data may be denied. The client IP address may be used to enforce security policies because a client IP address may be tightly coupled with the client identity, and the application services may obtain IP addresses when service request packets are received.

(36) However, port forwarding tunnel access techniques may replace the IP address of the client system in a packet with an IP address associated with the port forwarding tunnel server. Thus, the information associated with the client (e.g., the client IP address) may be lost, and security policies that rely on a client IP address may not be enforceable or useful. Various techniques may be used to address the described shortcomings of port forwarding tunnels, such as dropping of IP based security requirements or allowing all connections through port forwarding tunnels, but these techniques may result in increased security risks. Another solution is the use of an VPN solution that is complex, as the VPN solution may require an IP-encapsulation VPN (clients and server), a gateway, a reserved IP range as client VPN IP addresses, and related routes. Further, the VPN solutions may not address a problem of overlapping IP addresses. Overlapping IP may mean that the same IP address may be used by multiple clients, which may occur when multiple administrative domains issue IP addresses from similar ranges, especially from private IP address blocks. As such, a server (e.g., an application service) may not be able to differentiate clients based only on IP addresses, which may result in failure of client IP address based security policies.

(37) Techniques described herein support secure access to services via port forwarding tunneling techniques and using client IP based security policies. Specifically, an application server host may support virtual IP management and lookup. When a new client is provisioned for access to the server, the client is assigned a virtual IP address (e.g., a loopback IP address) that is maintained in a mapping that supports virtual IP address lookup using a client IP address and a client network identifier. The virtual IP is used to enforce security policies by the application server. A client may request access using a tunnel endpoint (e.g., a remote service server), and during connection setup, the virtual IP lookup is performed using the client IP address and the client network identifier (associated with the tunnel). From a tunnel service, the end service (e.g., application server) is connected using the corresponding virtual IP address. Thus, using these techniques, services provided by the computing system **105** (e.g., the server **125**), the DMS **110**, or both may be secured using client IP addresses and using port forwarding tunnel techniques to support secure access by

client systems (e.g., computing devices **115**).

(38) It should be appreciated by a person skilled in the art that one or more aspects of the disclosure may be implemented in the communicating environment **100** to additionally or alternatively solve other problems than those described herein. Further, aspects of the disclosure may provide technical improvements to “conventional” systems or processes as described herein. However, the description and appended drawings only include example technical improvements resulting from implementing aspects of the disclosure, and accordingly do not represent all of the technical improvements provided within the scope of the claims.

(39) FIG. 2 illustrates an example of a computing environment **200** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The computing environment includes a server network **205** and a client network **210**. The computing environment **200** may implement or may be implemented by aspects of computing environment **100** of FIG. 1. The server network **205** may support various computing devices that run host environments, such as application server host **215**, which may be an example of a virtual machine as described with respect to FIG. 1. The application server host **215** may support various services, such as an application server **225**. In some examples, the application server host **215** or the application server **225** is an example of DMS **110** of FIG. 1 and supports access to services by a client system, such as computing device **115** of FIG. 1

(40) The client network **210** may be an example of a local area network (LAN), virtual private network (VPN), or the like, that supports various services and clients, such as client hosts **250**. Each client host **250** may be associated with a client IP address. The client hosts **250** may access services of the application server **225** at a separate network using a port forwarding tunnel **265**, which may be implemented using SSH and SSF techniques, for example.

(41) Before a client host accesses services by the application server **225**, the client host **250** is provisioned with a virtual IP address at the application server host **215**. The application server host **215** may be configured with a range of virtual IP addresses from which to assign a virtual IP address to a client host **250** during client provisioning. In some examples, the range of virtual IP addresses includes a range of loopback IP addresses. The virtual IP address may not overlap with network addresses that are accessible via the application server host **215**. Loopback IP address blocks are defined in both IPv4 and IPv6, may be supported by most operating systems, and may be looped through a network interface card locally (e.g., the network interface **165** of FIG. 1). Depending on a size of a deployment and a number of potential client hosts **250**, a portion of the subnet of loopback IP addresses may be configured for assignment to hosts. When one client network (e.g., client network **210**) is accessing the application server host **215**, an entire or a portion of a subnet may be available for client hosts of the client network (e.g., client hosts **250** of client network **210**) the application server host **215**, each client network may be assigned a respective block of virtual IP addresses. For example, the client network **210** may be allocated the block 127.128.1.0/24, and virtual IP addresses of the client hosts **250** may be 127.128.1.1 and 127.12.1.2.

(42) When the client host **250** is provisioned, the application server host **215** may store an association of the IP address of the client host **250**, an identifier of the client network **210** (e.g., a client network name), and the newly assigned virtual IP address in the virtual IP mapping **230**. The virtual IP mapping **230** may be an example of a configuration file that is accessed by the application server host **215**, and the lines of the file may have the following format, but it should be understood that other formats are contemplated within the scope of the present disclosure: client network name 1, HostIP 1, Virtual IP 1 client network name 2, HostIP 2, Virtual IP 2

(43) This technique may be applicable when the deployment is relatively static (e.g., infrequent client host additions or deletions). When there are frequent host additions/deletions a virtual IP manager may be more suitable. In this implementation, a dedicated process may be created, and the dedicated process may listen to a TCP port. The TCP port may be used by the application server

host **215** for client host registration and deregistration. In such cases, the virtual IP mapping **230** may be maintained in a database.

(44) After a client host **250** is provisioned, the client host **250** may access the application server **225** via the tunnel **265**. For example, the client host **250-a** transmits a connection request to the application server host **215** via the port forwarding tunnel **265**. The connection request may include an indication of the client IP address for the client host **250-a**. In response to receiving the request, the application server host **215** may look up or identify a virtual IP address associated with the client host **250-a** using the client IP address and the client network name that is associated with the port forwarding tunnel **265** from which the request is received. The virtual IP address identification may be performed based on the implementation of the virtual IP mapping **230**. If the virtual IP mapping **230** is a lookup file, then the application server host **215** may access the file. If the virtual IP mapping **230** is implemented as a virtual IP manager, a lookup API request is transmitted by the application server host **215**, and a response to the API request may include an indication of the virtual IP address. The API request may include an indication of the client IP address and the client network identifier.

(45) Before the client host **250** is able to access the application server **225**, the tunnel **265** is to be deployed for connection between the client network **210** and the server network **205**. A tunnel endpoint **245-a** may be configured at the client hosts **250**. The tunnel endpoint **245-a** may be hosted by a tunnel endpoint host **240**. Further, a tunnel endpoint **245-a** may be a configured application server host **215** that hosts the application server **225** and tunnel forwarding rules may be provisioned to support client hosts **250** connecting to the application server **225**. The client hosts **250** are configured to use the IP address of the client tunnel endpoint **245-b** (e.g., an IP address of port **260**) to connect to the remote application server **225**. In some examples, the client hosts **250** are configured with the IP addresses using configuration files. As described herein, virtual IP addresses are allocated for the client hosts **250**, and the client host real IP address to virtual IP address mapping is saved as a configuration file on the application server host **215** or using a dedicated resolving server. On the application server host **215**, the client IP security is provisioned based on the corresponding virtual IPs. For example, for NFS, clients in the source IP block 127.128.1.0/24 may be permitted to read and write files in directory x/y/z.

(46) After the tunnel endpoints **245** and client hosts **250** are configured, the client side tunnel endpoint **245-b** (e.g., tunnel client) listens on the tunnel forwarding port configured for the remote application service. In some examples, the port may be the same as the listening port on the application server **225**. Client hosts **250** may transmit requests to the tunnel **265** at the client side using the port at which the tunnel endpoint **245-b** is listening. One tunnel endpoint **245-b** may support connection of multiple client hosts concurrently. As illustrated, both client hosts **250-a** and **250-b** are connected at the same time.

(47) For each client session towards the application server **225**, the tunnel client (e.g., tunnel endpoint **245-b**), forms a new connection to the tunnel endpoint **245-a** at the remote server (e.g., tunnel server). In the process, tunnel client (e.g., the tunnel endpoint **245-b**) may inform the tunnel server (e.g., the tunnel endpoint **245-a**) of the virtual IP address of the client host **250**. This information is known by the tunnel client because the client host **250** may connect to the tunnel client directly. For each client session, the tunnel server (e.g., the tunnel endpoint **245-a**) may perform a lookup to resolve the corresponding virtual IP of the client host **250**. For each session, the tunnel server may open a new socket, perform a socket bind using the identified virtual IP address of the client host, and initiate a connection to the application server **225** at the link local IP (e.g., 127.0.0.1 or the equivalent in IPv6). The application server **225** may see that the client connection is from a virtual IP (e.g., a loopback IP) and performs security checks based on the virtual IP. Based on the security policy, the service connection request is allowed or denied.

(48) As illustrated in FIG. 2, the connection request and/or other communications may be in the form of packets **270** (e.g., service packet for the connection requests). Communications from the

client host **250-a** may be in the form of packets **270-a**, and communications from the client host **250-b** may be in the form of packets **270-b**. A packet transmitted by a client host **250** may initially include the IP address of the client host **250**, but the IP address may be replaced by a source IP address of the port forwarding tunnel **265**. The tunnel endpoint **245-a**, upon receiving the packet with the source IP address, may identify the virtual IP address corresponding to the client host **250** (e.g., based on the client IP address received from the tunnel endpoint **245-**) and the identifier for the client network **210**. The tunnel endpoint **245-a** may replace the source IP address of the packet with the identified virtual IP address for the client host and forward the packet to the application server **225**. As such, the application server **225** (or the application server host **215**) may enforce security policies **235** based on the IP address that is included in the received packets.

(49) FIG. 3 illustrates an example of a computing environment **300** for multiple client networks that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The computing environment **300** may implement aspects of computing environment **100** of FIG. 1 and computing environment **200** of FIG. 2. For example, the computing environment **300** includes a server network **305**, which may be an example of server network **205** of FIG. 2. The computing environment **300** also includes client networks **310**, which may be examples of client network **210** of FIG. 2. The server network **305** includes an application server host **315**, which maintains a virtual IP mapping **330** that is used to support secure access to an application server **325** by client hosts **350** via port forwarding tunnels **365**.

(50) As illustrated in FIG. 3, client hosts **350** of two different client networks **310-a** and **310-b** may be configured to access the application server **325** via respective port forwarding tunnels **365-a** and **365-b**. Each client network **310** may correspond to a different organization, team, location, or the like. Each client host **350** may be associated with a respective client IP address **370**. The techniques described herein support security policy enforcement for access to the application server **325** when multiple client hosts **350** have overlapping IP addresses. Overlapping IP addresses means the same IP address may be used by multiple clients. This may occur when there are multiple administrative domains (e.g., client networks **310**), each issuing IPs from similar ranges (e.g., especially from private IP address blocks). It is possible that these domains may share the same application service (e.g., application server **325**). Techniques described herein support policy enforcement when client hosts share a same IP address **370**.

(51) As illustrated in FIG. 3, client host **350-a** of client network **310-a** has a IP address **370-a** “IP address A.” Client host **350-c** of client network **310-b** has an IP address **370-c** “IP address A.” As such, the IP address **370-a** and the IP address **370-c** are the same. Using the techniques described herein, the application server **325** is able to implement IP based security policies while being able to differentiate between the client hosts **350-a** and client host **350-c**. Each of the client hosts **350** is assigned a respective virtual IP address by the application server host **315**, and the application server hosts **315** maintains the virtual IP mapping **330**.

(52) When a client host **350** connects to the application server **325** via the tunnel, the application server host **315** identifies the virtual IP address that is assigned to the client host **350** using the virtual IP mapping **330**. The application server host **315** may identify the virtual IP address using the respective client IP address **370** and a client network identifier corresponding to the client network **310** in which the client host **350** operates. Thus, when two client IP addresses are the same (e.g., client IP address **370-a** and client IP address **370-c**), the application server host **315** uses the corresponding client network name to differentiate between the client hosts **350** and the client IP addresses **370-c**. In some examples, the client network name may be determined based on the tunnel **365** via which a connection request is received. For example, as the client host **350-a** sends a connection request via the tunnel endpoint **340-a** and the tunnel **365-a**, the application server host **315** (e.g., a tunnel endpoint **345**) is able to identify the corresponding client network name of the client network **310-a** based on the use of the tunnel **365-a** that is used to transmit the connection request. Similarly, as the client host **350-c** sends a connection request via the tunnel endpoint **340-b**

and the tunnel **365-b**, the application server host **315** (e.g., a tunnel endpoint **345**) is able to identify the corresponding client network name of the client network **310-b** based on the use of the tunnel **365-b** that is used to transmit the connection request. As such, techniques described herein support security policies for application servers **325** using port forwarding tunnels **365** when client hosts **350** of different domains (e.g., client networks **310**) may share overlapping IP addresses.

(53) FIG. **4** illustrates an example of a process flow **400** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The process flow **400** includes an application server host **405** and a client host **410**, which may be examples of the corresponding devices described with respect to FIGS. **1** through **3**. In some examples, some signaling or procedure of the process flow **400** may occur in different orders than shown. Additionally, or alternatively, some additional procedures of signaling may occur, or some signaling or procedures may not occur.

(54) At **415**, the application server host **405** may receive, via a port forwarding tunnel, a connection request that includes an indication of a client IP address for the client host **410**. The connection request may be in the form of one or more service request packets that include an IP address associated with the port forwarding tunnel. In such cases, a tunneling endpoint (e.g., a tunnel client) may inform the application server host **405** (e.g., a tunnel server) of the client IP address associated with the client host **410** that transmitted the request).

(55) At **420**, the application server host **405** (e.g., a tunnel server) may identify a client IP address associated with the client host **410** and a client network identifier associated with the client host **410**. The client network identifier may be identified based on the port forwarding tunnel via which the request is received. As such, the tunnel server may identify the client network identifier.

(56) At **425**, the application server host **405** (e.g., a tunnel server) may identify, based at least in part on the client IP address and using a mapping maintained for a plurality of client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. In some examples, the application server host **405** may identify the virtual IP address by identifying the virtual IP address from a file including the mapping and maintained at the application server host.

Additionally, or alternatively, the application server host **405** may transmit, via an API, a request for the virtual IP address, wherein the request includes the client IP address and the client network identifier and receive, in response to the request, an indication of the virtual IP address.

(57) At **430**, the application server host **405** may open a connection between an application server hosted by the application server host **405** and the client host **410** using the identified virtual IP address. In some examples, the application server host **405** (e.g., a tunnel server) may replace a source IP address (in the packet) associated with the port forwarding tunnel with the virtual IP address associated with the client host **410** and forward the packet to the application server hosted by the application server host **405**.

(58) At **435**, the application server of the application server host **405** and the client host **410** may communicate via the opened connection.

(59) FIG. **5** shows a block diagram **500** of a system **505** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. In some examples, the system **505** may be an example of aspects of one or more components described with respect to FIG. **1**, such as a DMS **110** or a server **125**. The system **505** may include an input interface **510**, an output interface **515**, and a connection manager **520**. The system **505** may also include one or more processors. Each of these components may be in communication with one another (e.g., via one or more communications links, communications interfaces, or any combination thereof).

(60) The input interface **510** may manage input signals for the system **505**. For example, the input interface **510** may receive input signaling (e.g., messages, packets, data, instructions, commands, or any other form of encoded information) from other systems or devices. The input interface **510** may send signaling corresponding to (e.g., representative of or otherwise based on) such input

signaling to other components of the system **505** for processing. For example, the input interface **510** may transmit such corresponding signaling to the connection manager **520** to support IP based security over port forwarding tunnels. In some cases, the input interface **510** may be a component of a network interface **710** as described with reference to FIG. 7.

(61) The output interface **515** may manage output signaling for the system **505**. For example, the output interface **515** may receive signaling from other components of the system **505**, such as the connection manager **520**, and may transmit such output signaling corresponding to (e.g., representative of or otherwise based on) such signaling to other systems or devices. In some cases, the output interface **515** may be a component of a network interface **710** as described with reference to FIG. 7

(62) The connection manager **520** may include a connection interface **525**, a virtual IP identification component **530**, a connection opening component **535**, or any combination thereof. In some examples, the connection manager **520**, or various components thereof, may be configured to perform various operations (e.g., receiving, monitoring, transmitting) using or otherwise in cooperation with the input interface **510**, the output interface **515**, or both. For example, the connection manager **520** may receive information from the input interface **510**, send information to the output interface **515**, or be integrated in combination with the input interface **510**, the output interface **515**, or both to receive information, transmit information, or perform various other operations as described herein.

(63) The connection interface **525** may be configured as or otherwise support a means for receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The virtual IP identification component **530** may be configured as or otherwise support a means for identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The connection opening component **535** may be configured as or otherwise support a means for opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(64) FIG. 6 shows a block diagram **600** of a connection manager **620** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The connection manager **620** may be an example of aspects of a connection manager **520** as described herein. The connection manager **620**, or various components thereof, may be an example of means for performing various aspects of IP based security over port forwarding tunnels as described herein. For example, the connection manager **620** may include a connection interface **625**, a virtual IP identification component **630**, a connection opening component **635**, a file access component **640**, an API component **645**, an API component **650**, a packet interface **655**, an IP address replacement component **660**, a packet forwarding component **665**, or any combination thereof. Each of these components may communicate, directly or indirectly, with one another (e.g., via one or more buses communications links, communications interfaces, or any combination thereof).

(65) The connection interface **625** may be configured as or otherwise support a means for receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The virtual IP identification component **630** may be configured as or otherwise support a means for identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The connection opening component **635** may be configured as or otherwise support a means for opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(66) In some examples, to support identifying the virtual IP address, the file access component **640** may be configured as or otherwise support a means for identifying, based on the client IP address and the client network identifier, the virtual IP address from a file including the mapping and maintained at the application server host.

(67) In some examples, to support identifying the virtual IP address, the API component **645** may be configured as or otherwise support a means for transmitting, via an application programming interface (API), a request for the virtual IP address, where the request includes the client IP address and the client network identifier. In some examples, to support identifying the virtual IP address, the API component **650** may be configured as or otherwise support a means for receiving, in response to the request, an indication of the virtual IP address.

(68) In some examples, the connection interface **625** may be configured as or otherwise support a means for receiving, at an application server host and via a second port forwarding tunnel, a second connection request that includes an indication of a second client IP address for a second client host, where the second client IP address has a same value as the client IP address. In some examples, the virtual IP identification component **630** may be configured as or otherwise support a means for identifying, based on the second client IP address and using the mapping, a second virtual IP address that is associated with the second client IP address and a second client network identifier associated with the second port forwarding tunnel from which the connection request is received, where the second client network identifier is used to differentiate between the second client IP address and the client IP address that have the same value. In some examples, the connection opening component **635** may be configured as or otherwise support a means for opening a second connection between the application server hosted and the second client host using the identified second virtual IP address.

(69) In some examples, to support receiving the connection request, the connection interface **625** may be configured as or otherwise support a means for receiving the connection request via a tunnel endpoint for the port forwarding tunnel, where the tunnel endpoint is provisioned for access by a set of multiple client hosts associated with the client network identifier.

(70) In some examples, the set of multiple client hosts are associated with a set of multiple loopback IP addresses in the mapping. In some examples, the application server host is configured to enforce security policies based on the identified virtual IP address.

(71) In some examples, the packet interface **655** may be configured as or otherwise support a means for receiving, via the connection, a packet that is to be received at the application server, where the packet includes an IP address associated with the port forwarding tunnel. In some examples, the IP address replacement component **660** may be configured as or otherwise support a means for replacing the IP address associated with the port forwarding tunnel with the virtual IP address associated with the client host based on the connection via which the packet is received. In some examples, the packet forwarding component **665** may be configured as or otherwise support a means for forwarding the packet to the application server.

(72) FIG. 7 shows a block diagram **700** of a system **705** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The system **705** may be an example of or include the components of a system **505** as described herein. The system **705** may include components for data management, such as a connection manager **720**, a network interface **710**, memory **725**, processor **730**, and storage **735**. These components may be in electronic communication or otherwise coupled with each other (e.g., operatively, communicatively, functionally, electronically, electrically; via one or more buses, communications links, communications interfaces, or any combination thereof). Additionally, the components of the system **505** may comprise corresponding physical components or may be implemented as corresponding virtual components (e.g., components of one or more virtual machines). In some examples, the system **705** may be an example of aspects of one or more components described with reference to FIG. 1, such as a DMS **110** or server **125**, or one or application servers or client

networks as described with respect to FIGS. 2 through 4.

(73) The network interface **710** may enable the system **705** to exchange information (e.g., input information **745**, output information, **750**, or both) with other systems or devices (not shown). For example, the network interface **710** may include one or more wireless network interfaces, one or more wired network interfaces, or any combination thereof. In some examples, the network interface **710** may be an example of aspects of one or more components described with reference to FIG. 1, such as one or more network interfaces **165** or similar components of other systems such as server **125** or application server hosts or client network as described with respect to FIGS. 2 through 4.

(74) Memory **725** may include RAM, ROM, or both. The memory **725** may store computer-readable, computer-executable software including instructions that, when executed, cause the processor **730** to perform various functions described herein. In some cases, the memory **725** may contain, among other things, a basic input/output system (BIOS), which may control basic hardware or software operation such as the interaction with peripheral components or devices. In some cases, the memory **725** may be an example of aspects of one or more components described with reference to FIG. 1, such as one or more memories **175**, or one or more aspects of components of FIGS. 2 through 4, such as aspects of the application server host or the client network.

(75) The processor **730** may include an intelligent hardware device, (e.g., a general-purpose processor, a digital signal processor (DSP), a CPU, a microcontroller, an application-specific integrated circuit (ASIC), a field programmable gate array (FPGA), a programmable logic device, a discrete gate or transistor logic component, a discrete hardware component, or any combination thereof). The processor **730** may be configured to execute computer-readable instructions stored in the memory **725** to perform various functions (e.g., functions or tasks supporting IP based security over port forwarding tunnels). Through a single processor **730** is depicted in the example of FIG. 7, it should be understood that the system **705** may include any quantity of one or more of processors **730** and that a group of processors **730** may collectively perform one or more functions ascribed herein to a processor, such as the processor **730**. In some cases, the processor **730** may be an example of aspects of one or more components described with reference to FIG. 1, such as one or more processors **170**, or one or more aspects of components of FIGS. 2 through 4, such as aspects of the application server host or the client network.

(76) Storage **735** may be configured to store data that is generated, processed, stored, or otherwise used by the system **705**. In some cases, the storage **735** may include one or more HDDs, one or more SSDs, or both. In some examples, the storage **735** may be an example of a single database, a distributed database, multiple distributed databases, a data store, a data lake, or an emergency backup database. In some examples, the storage **735** may be an example of one or more components described with reference to FIG. 1, such as one or more network disks **180** or one or more aspects of components of FIGS. 2 through 4, such as aspects of the application server host or the client network.

(77) The connection manager **720** may be configured as or otherwise support a means for receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The connection manager **720** may be configured as or otherwise support a means for identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The connection manager **720** may be configured as or otherwise support a means for opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(78) By including or configuring the connection manager **720** in accordance with examples as described herein, the system **705** may support techniques for secure access to computing resources, such as resources of an application server using virtual IP based security policies. The techniques

may be secured via a virtual IP mapping and may support secure port forwarding tunneling techniques with or more client domains (e.g., client networks).

(79) FIG. 8 shows a flowchart illustrating a method **800** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The operations of the method **800** may be implemented by a server host or its components as described herein. For example, the operations of the method **800** may be performed by a DMS or server host as described with reference to FIGS. 1 through 7. In some examples, a DMS or server host may execute a set of instructions to control the functional elements of the server host to perform the described functions. Additionally, or alternatively, the DMS or server host may perform aspects of the described functions using special-purpose hardware.

(80) At **805**, the method may include receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The operations of **805** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **805** may be performed by a connection interface **625** as described with reference to FIG. 6.

(81) At **810**, the method may include identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The operations of **810** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **810** may be performed by a virtual IP identification component **630** as described with reference to FIG. 6.

(82) At **815**, the method may include opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address. The operations of **815** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **815** may be performed by a connection opening component **635** as described with reference to FIG. 6.

(83) FIG. 9 shows a flowchart illustrating a method **900** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The operations of the method **900** may be implemented by a server host or its components as described herein. For example, the operations of the method **900** may be performed by a server host as described with reference to FIGS. 1 through 7. In some examples, a DMS or server host may execute a set of instructions to control the functional elements of the server host to perform the described functions. Additionally, or alternatively, the DMS or server host may perform aspects of the described functions using special-purpose hardware.

(84) At **905**, the method may include receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The operations of **905** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **905** may be performed by a connection interface **625** as described with reference to FIG. 6.

(85) At **910**, the method may include identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The operations of **910** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **910** may be performed by a virtual IP identification component **630** as described with reference to FIG. 6.

(86) At **915**, the method may include identifying, based on the client IP address and the client network identifier, the virtual IP address from a file including the mapping and maintained at the application server host. The operations of **915** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **915** may be performed by a file access component **640** as described with reference to FIG. 6.

(87) At **920**, the method may include opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address. The operations of **920** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **920** may be performed by a connection opening component **635** as described with reference to FIG. 6.

(88) FIG. 10 shows a flowchart illustrating a method **1000** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The operations of the method **1000** may be implemented by a server host or its components as described herein. For example, the operations of the method **1000** may be performed by a server host as described with reference to FIGS. 1 through 7. In some examples, a DMS or server host may execute a set of instructions to control the functional elements of the server host to perform the described functions. Additionally, or alternatively, the DMS or server host may perform aspects of the described functions using special-purpose hardware.

(89) At **1005**, the method may include receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The operations of **1005** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1005** may be performed by a connection interface **625** as described with reference to FIG. 6.

(90) At **1010**, the method may include identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The operations of **1010** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1010** may be performed by a virtual IP identification component **630** as described with reference to FIG. 6.

(91) At **1015**, the method may include transmitting, via an application programming interface (API), a request for the virtual IP address, where the request includes the client IP address and the client network identifier. The operations of **1015** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1015** may be performed by an API component **645** as described with reference to FIG. 6.

(92) At **1020**, the method may include receiving, in response to the request, an indication of the virtual IP address. The operations of **1020** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1020** may be performed by an API component **650** as described with reference to FIG. 6.

(93) At **1025**, the method may include opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address. The operations of **1025** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1025** may be performed by a connection opening component **635** as described with reference to FIG. 6.

(94) FIG. 11 shows a flowchart illustrating a method **1100** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The operations of the method **1100** may be implemented by a server host or its components as described herein. For example, the operations of the method **1100** may be performed by a server host as described with reference to FIGS. 1 through 7. In some examples, a DMS or server host may execute a set of instructions to control the functional elements of the server host to perform the described functions. Additionally, or alternatively, the DMS or server host may perform aspects of the described functions using special-purpose hardware.

(95) At **1105**, the method may include receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The operations of **1105** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1105** may be performed by a connection

interface **625** as described with reference to FIG. **6**.

(96) At **1110**, the method may include identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The operations of **1110** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1110** may be performed by a virtual IP identification component **630** as described with reference to FIG. **6**.

(97) At **1115**, the method may include opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address. The operations of **1115** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1115** may be performed by a connection opening component **635** as described with reference to FIG. **6**.

(98) At **1120**, the method may include receiving, at an application server host and via a second port forwarding tunnel, a second connection request that includes an indication of a second client IP address for a second client host, where the second client IP address has a same value as the client IP address. The operations of **1120** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1120** may be performed by a connection interface **625** as described with reference to FIG. **6**.

(99) At **1125**, the method may include identifying, based on the second client IP address and using the mapping, a second virtual IP address that is associated with the second client IP address and a second client network identifier associated with the second port forwarding tunnel from which the connection request is received, where the second client network identifier is used to differentiate between the second client IP address and the client IP address that have the same value. The operations of **1125** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1125** may be performed by a virtual IP identification component **630** as described with reference to FIG. **6**.

(100) At **1130**, the method may include opening a second connection between the application server hosted and the second client host using the identified second virtual IP address. The operations of **1130** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1130** may be performed by a connection opening component **635** as described with reference to FIG. **6**.

(101) FIG. **12** shows a flowchart illustrating a method **1200** that supports IP based security over port forwarding tunnels in accordance with aspects of the present disclosure. The operations of the method **1200** may be implemented by a server host or its components as described herein. For example, the operations of the method **1200** may be performed by a server host as described with reference to FIGS. **1** through **7**. In some examples, a DMS or server host may execute a set of instructions to control the functional elements of the server host to perform the described functions. Additionally, or alternatively, the DMS or server host may perform aspects of the described functions using special-purpose hardware.

(102) At **1205**, the method may include receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host. The operations of **1205** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1205** may be performed by a connection interface **625** as described with reference to FIG. **6**.

(103) At **1210**, the method may include receiving, via the connection, a packet that is to be received at the application server, where the packet includes an IP address associated with the port forwarding tunnel. The operations of **1210** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1210** may be performed by a packet interface **655** as described with reference to FIG. **6**.

(104) At **1215**, the method may include identifying, based on the client IP address and using a

mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received. The operations of **1215** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1215** may be performed by a virtual IP identification component **630** as described with reference to FIG. 6.

(105) At **1220**, the method may include replacing the IP address associated with the port forwarding tunnel with the virtual IP address associated with the client host based on the connection via which the packet is received. The operations of **1220** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1220** may be performed by an IP address replacement component **660** as described with reference to FIG. 6.

(106) At **1225**, the method may include opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address. The operations of **1225** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1225** may be performed by a connection opening component **635** as described with reference to FIG. 6.

(107) At **1230**, the method may include forwarding the packet to the application server. The operations of **1230** may be performed in accordance with examples as disclosed herein. In some examples, aspects of the operations of **1230** may be performed by a packet forwarding component **665** as described with reference to FIG. 6.

(108) A method is described. The method may include receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host, identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received, and opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(109) An apparatus is described. The apparatus may include a processor, memory coupled with the processor, and instructions stored in the memory. The instructions may be executable by the processor to cause the apparatus to receive, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host, identify, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received, and open a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(110) Another apparatus is described. The apparatus may include means for receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host, means for identifying, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is received, and means for opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(111) A non-transitory computer-readable medium storing code is described. The code may include instructions executable by a processor to receive, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client IP (IP) address for a client host, identify, based on the client IP address and using a mapping maintained for a set of multiple client hosts, a virtual IP address that is associated with the client IP address and a client network identifier associated with the port forwarding tunnel from which the connection request is

received, and open a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

(112) In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, identifying the virtual IP address may include operations, features, means, or instructions for identifying, based on the client IP address and the client network identifier, the virtual IP address from a file including the mapping and maintained at the application server host.

(113) In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, identifying the virtual IP address may include operations, features, means, or instructions for transmitting, via an application programming interface (API), a request for the virtual IP address, where the request includes the client IP address and the client network identifier and receiving, in response to the request, an indication of the virtual IP address.

(114) Some examples of the method, apparatuses, and non-transitory computer-readable medium described herein may further include operations, features, means, or instructions for receiving, at an application server host and via a second port forwarding tunnel, a second connection request that includes an indication of a second client IP address for a second client host, where the second client IP address may have a same value as the client IP address, identifying, based on the second client IP address and using the mapping, a second virtual IP address that may be associated with the second client IP address and a second client network identifier associated with the second port forwarding tunnel from which the connection request may be received, where the second client network identifier may be used to differentiate between the second client IP address and the client IP address that may have the same value, and opening a second connection between the application server hosted and the second client host using the identified second virtual IP address.

(115) In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, receiving the connection request may include operations, features, means, or instructions for receiving the connection request via a tunnel endpoint for the port forwarding tunnel, where the tunnel endpoint may be provisioned for access by a set of multiple client hosts associated with the client network identifier.

(116) In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, the set of multiple client hosts may be associated with a set of multiple loopback IP addresses in the mapping.

(117) In some examples of the method, apparatuses, and non-transitory computer-readable medium described herein, the application server host may be configured to enforce security policies based on the identified virtual IP address.

(118) Some examples of the method, apparatuses, and non-transitory computer-readable medium described herein may further include operations, features, means, or instructions for receiving, via the connection, a packet that may be received at the application server, where the packet includes an IP address associated with the port forwarding tunnel, replacing the IP address associated with the port forwarding tunnel with the virtual IP address associated with the client host based on the connection via which the packet may be received, and forwarding the packet to the application server.

(119) It should be noted that the methods described above describe possible implementations, and that the operations and the steps may be rearranged or otherwise modified and that other implementations are possible. Further, aspects from two or more of the methods may be combined.

(120) The description set forth herein, in connection with the appended drawings, describes example configurations and does not represent all the examples that may be implemented or that are within the scope of the claims. The term “exemplary” used herein means “serving as an example, instance, or illustration,” and not “preferred” or “advantageous over other examples.” The detailed description includes specific details for the purpose of providing an understanding of the described techniques. These techniques, however, may be practiced without these specific details. In some instances, well-known structures and devices are shown in block diagram form in order to

avoid obscuring the concepts of the described examples.

(121) In the appended figures, similar components or features may have the same reference label. Further, various components of the same type may be distinguished by following the reference label by a dash and a second label that distinguishes among the similar components. If just the first reference label is used in the specification, the description is applicable to any one of the similar components having the same first reference label irrespective of the second reference label.

(122) Information and signals described herein may be represented using any of a variety of different technologies and techniques. For example, data, instructions, commands, information, signals, bits, symbols, and chips that may be referenced throughout the above description may be represented by voltages, currents, electromagnetic waves, magnetic fields or particles, optical fields or particles, or any combination thereof.

(123) The various illustrative blocks and modules described in connection with the disclosure herein may be implemented or performed with a general-purpose processor, a DSP, an ASIC, an FPGA or other programmable logic device, discrete gate or transistor logic, discrete hardware components, or any combination thereof designed to perform the functions described herein. A general-purpose processor may be a microprocessor, but in the alternative, the processor may be any conventional processor, controller, microcontroller, or state machine. A processor may also be implemented as a combination of computing devices (e.g., a combination of a DSP and a microprocessor, multiple microprocessors, one or more microprocessors in conjunction with a DSP core, or any other such configuration).

(124) The functions described herein may be implemented in hardware, software executed by a processor, firmware, or any combination thereof. If implemented in software executed by a processor, the functions may be stored on or transmitted over as one or more instructions or code on a computer-readable medium. Other examples and implementations are within the scope of the disclosure and appended claims. For example, due to the nature of software, functions described above can be implemented using software executed by a processor, hardware, firmware, hardwiring, or combinations of any of these. Features implementing functions may also be physically located at various positions, including being distributed such that portions of functions are implemented at different physical locations. Also, as used herein, including in the claims, “or” as used in a list of items (for example, a list of items prefaced by a phrase such as “at least one of” or “one or more of”) indicates an inclusive list such that, for example, a list of at least one of A, B, or C means A or B or C or AB or AC or BC or ABC (i.e., A and B and C). Also, as used herein, the phrase “based on” shall not be construed as a reference to a closed set of conditions. For example, an exemplary step that is described as “based on condition A” may be based on both a condition A and a condition B without departing from the scope of the present disclosure. In other words, as used herein, the phrase “based on” shall be construed in the same manner as the phrase “based at least in part on.”

(125) Computer-readable media includes both non-transitory computer storage media and communication media including any medium that facilitates transfer of a computer program from one place to another. A non-transitory storage medium may be any available medium that can be accessed by a general purpose or special purpose computer. By way of example, and not limitation, non-transitory computer-readable media can comprise RAM, ROM, electrically erasable programmable ROM (EEPROM), compact disk (CD) ROM or other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other non-transitory medium that can be used to carry or store desired program code means in the form of instructions or data structures and that can be accessed by a general-purpose or special-purpose computer, or a general-purpose or special-purpose processor. Also, any connection is properly termed a computer-readable medium. For example, if the software is transmitted from a website, server, or other remote source using a coaxial cable, fiber optic cable, twisted pair, digital subscriber line (DSL), or wireless technologies such as infrared, radio, and microwave, then the coaxial cable, fiber optic cable, twisted pair, DSL,

or wireless technologies such as infrared, radio, and microwave are included in the definition of medium. Disk and disc, as used herein, include CD, laser disc, optical disc, digital versatile disc (DVD), floppy disk and Blu-ray disc where disks usually reproduce data magnetically, while discs reproduce data optically with lasers. Combinations of the above are also included within the scope of computer-readable media.

(126) The description herein is provided to enable a person skilled in the art to make or use the disclosure. Various modifications to the disclosure will be readily apparent to those skilled in the art, and the generic principles defined herein may be applied to other variations without departing from the scope of the disclosure. Thus, the disclosure is not limited to the examples and designs described herein, but is to be accorded the broadest scope consistent with the principles and novel features disclosed herein.

Claims

1. A method for data management comprising: receiving, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client internet protocol (IP) address for a client host; identifying, based at least in part on the client IP address and using a mapping maintained for a plurality of client hosts, a virtual IP address that is assigned to both the client IP address and a client network identifier, wherein the client network identifier is mapped to the port forwarding tunnel from which the connection request is received; and opening a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.
2. The method of claim 1, wherein identifying the virtual IP address comprises: identifying, based at least in part on the client IP address and the client network identifier, the virtual IP address from a file including the mapping and maintained at the application server host.
3. The method of claim 1, wherein identifying the virtual IP address comprises: transmitting, via an application programming interface (API), a request for the virtual IP address, wherein the request includes the client IP address and the client network identifier; and receiving, in response to the request, an indication of the virtual IP address.
4. The method of claim 1, further comprising: receiving, at the application server host and via a second port forwarding tunnel, a second connection request that includes an indication of a second client IP address for a second client host, wherein the second client IP address has a same value as the client IP address; identifying, based at least in part on the second client IP address and using the mapping, a second virtual IP address that is assigned to both the second client IP address and a second client network identifier, wherein the client network identifier is mapped to the second port forwarding tunnel from which the connection request is received, wherein the second client network identifier is used to differentiate between the second client IP address and the client IP address that have the same value; and opening a second connection between the application server hosted and the second client host using the identified second virtual IP address.
5. The method of claim 1, wherein receiving the connection request comprises: receiving the connection request via a tunnel endpoint for the port forwarding tunnel, wherein the tunnel endpoint is provisioned for access by a second plurality of client hosts associated with the client network identifier.
6. The method of claim 1, wherein the plurality of client hosts are associated with a plurality of loopback IP addresses in the mapping.
7. The method of claim 1, wherein the application server host is configured to enforce security policies based at least in part on the identified virtual IP address.
8. The method of claim 1, further comprising: receiving, via the connection, a packet that is to be received at the application server, wherein the packet includes an IP address associated with the port forwarding tunnel; replacing the IP address associated with the port forwarding tunnel with the

virtual IP address associated with the client host based at least in part on the connection via which the packet is received; and forwarding the packet to the application server.

9. An apparatus, comprising: a processor; memory coupled with the processor; and instructions stored in the memory and executable by the processor to cause the apparatus to: receive, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client internet protocol (IP) address for a client host; identify, based at least in part on the client IP address and using a mapping maintained for a plurality of client hosts, a virtual IP address that is assigned to both the client IP address and a client network identifier, wherein the client network identifier is mapped to the port forwarding tunnel from which the connection request is received; and open a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

10. The apparatus of claim 9, wherein the instructions to identify the virtual IP address are executable by the processor to cause the apparatus to: identify, based at least in part on the client IP address and the client network identifier, the virtual IP address from a file including the mapping and maintained at the application server host.

11. The apparatus of claim 9, wherein the instructions to identify the virtual IP address are executable by the processor to cause the apparatus to: transmit, via an application programming interface (API), a request for the virtual IP address, wherein the request includes the client IP address and the client network identifier; and receive, in response to the request, an indication of the virtual IP address.

12. The apparatus of claim 9, wherein the instructions are further executable by the processor to cause the apparatus to: receive, at the application server host and via a second port forwarding tunnel, a second connection request that includes an indication of a second client IP address for a second client host, wherein the second client IP address has a same value as the client IP address; identify, based at least in part on the second client IP address and using the mapping, a second virtual IP address that is assigned to both the second client IP address and a second client network identifier, wherein the client network identifier is mapped to the second port forwarding tunnel from which the connection request is received, wherein the second client network identifier is used to differentiate between the second client IP address and the client IP address that have the same value; and open a second connection between the application server hosted and the second client host using the identified second virtual IP address.

13. The apparatus of claim 9, wherein the instructions to receive the connection request are executable by the processor to cause the apparatus to: receive the connection request via a tunnel endpoint for the port forwarding tunnel, wherein the tunnel endpoint is provisioned for access by a second plurality of client hosts associated with the client network identifier.

14. The apparatus of claim 9, wherein the plurality of client hosts are associated with a plurality of loopback IP addresses in the mapping.

15. The apparatus of claim 9, wherein the application server host is configured to enforce security policies based at least in part on the identified virtual IP address.

16. The apparatus of claim 9, wherein the instructions are further executable by the processor to cause the apparatus to: receive, via the connection, a packet that is to be received at the application server, wherein the packet includes an IP address associated with the port forwarding tunnel; replace the IP address associated with the port forwarding tunnel with the virtual IP address associated with the client host based at least in part on the connection via which the packet is received; and forward the packet to the application server.

17. A non-transitory computer-readable medium storing code, the code comprising instructions executable by a processor to: receive, at an application server host and via a port forwarding tunnel, a connection request that includes an indication of a client internet protocol (IP) address for a client host; identify, based at least in part on the client IP address and using a mapping maintained for a plurality of client hosts, a virtual IP address that is assigned to both the client IP address and a

client network identifier, wherein the client network identifier is mapped to the port forwarding tunnel from which the connection request is received; and open a connection between an application server hosted by the application server host and the client host using the identified virtual IP address.

18. The non-transitory computer-readable medium of claim 17, wherein the instructions are further executable by the processor to: receive, at the application server host and via a second port forwarding tunnel, a second connection request that includes an indication of a second client IP address for a second client host, wherein the second client IP address has a same value as the client IP address; identify, based at least in part on the second client IP address and using the mapping, a second virtual IP address that is assigned to both the second client IP address and a second client network identifier, wherein the client network identifier is mapped to the second port forwarding tunnel from which the connection request is received, wherein the second client network identifier is used to differentiate between the second client IP address and the client IP address that have the same value; and open a second connection between the application server hosted and the second client host using the identified second virtual IP address.

19. The non-transitory computer-readable medium of claim 17, wherein the instructions to receive the connection request are executable by the processor to: receive the connection request via a tunnel endpoint for the port forwarding tunnel, wherein the tunnel endpoint is provisioned for access by a second plurality of client hosts associated with the client network identifier.

20. The non-transitory computer-readable medium of claim 17, wherein the plurality of client hosts are associated with a plurality of loopback IP addresses in the mapping.
