

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12418657
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Martinelli; Ilaria et al.

Guaranteed data compression

Abstract

A method of converting 10-bit pixel data (e.g. 10:10:10:2 data) into 8-bit pixel data involves converting the 10-bit values to 7-bits or 8-bits and generating error values for each of the converted values. Two of the 8-bit output channels comprise a combination of a converted 7-bit value and one of the bits from the fourth input channel. A third 8-bit output channel comprises the converted 8-bit value and the fourth 8-bit output channel comprises the error values. In various examples, the bits of the error values may be interleaved when they are packed into the fourth output channel.

Inventors: Martinelli; Ilaria (Watford, GB), Bond; Jeff (Barry, GB), Fenney; Simon (St. Albans, GB), Lacey; Peter Malcolm (Bukinghamshire, GB), Clark; Gregory (Hemel Hempstead, GB)

Applicant: Imagination Technologies Limited (Kings Langley, GB)

Family ID: 1000008818905

Assignee: Imagination Technologies Limited (Kings Langley, GB)

Appl. No.: 18/671950

Filed: May 22, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240314323 A1	Sep. 19, 2024

Publication Classification

Int. Cl.: H04N19/00 (20140101); G06T1/20 (20060101); H04N19/132 (20140101); H04N19/176 (20140101); H04N19/42 (20140101); H04N19/89 (20140101)

U.S. Cl.:

CPC H04N19/132 (20141101); G06T1/20 (20130101); H04N19/176 (20141101); H04N19/42 (20141101); H04N19/89 (20141101);

Field of Classification Search

CPC: H04N (19/00); G06T (1/00)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2673065	12/1953	Patterson	N/A	N/A
5073776	12/1990	Shigemori	N/A	N/A
5218563	12/1992	Juri et al.	N/A	N/A
5299208	12/1993	Blaum et al.	N/A	N/A
5673065	12/1996	DeLeeuw	N/A	N/A

5751359	12/1997	Suzuki et al.	N/A	N/A
5809200	12/1997	Nishimoto et al.	N/A	N/A
5859857	12/1998	Martinson et al.	N/A	N/A
6137912	12/1999	Kostrzewski et al.	N/A	N/A
6580755	12/2002	Morimoto et al.	N/A	N/A
6583887	12/2002	Clouthier et al.	N/A	N/A
6771727	12/2003	Haverinen	N/A	N/A
6771827	12/2003	Curry	N/A	N/A
7245663	12/2006	Schaar et al.	N/A	N/A
7920749	12/2010	Donovan	N/A	N/A
8428371	12/2012	Morovic et al.	N/A	N/A
9215505	12/2014	Pelerin	N/A	N/A
9275605	12/2015	Longhurst et al.	N/A	N/A
9451254	12/2015	Joshi et al.	N/A	N/A
9998142	12/2017	Gopal	N/A	N/A
10225569	12/2018	Okamoto et al.	N/A	N/A
10707895	12/2019	Fenney	N/A	N/A
10720940	12/2019	Fenney et al.	N/A	N/A
10812101	12/2019	Lacey et al.	N/A	N/A
10819367	12/2019	Fenney et al.	N/A	N/A
10895776	12/2020	Yang	N/A	N/A
10985776	12/2020	Fenney	N/A	N/A
11070227	12/2020	Fenney et al.	N/A	N/A
11258457	12/2021	Lacey et al.	N/A	N/A
11309907	12/2021	Fenney et al.	N/A	N/A
2003/0030575	12/2002	Frachtenberg et al.	N/A	N/A
2003/0090397	12/2002	Rasmussen	N/A	N/A
2004/0042292	12/2003	Sakata et al.	N/A	N/A
2004/0189679	12/2003	Ito	345/690	G09G 3/20
2005/0200631	12/2004	Pan et al.	N/A	N/A
2006/0115166	12/2005	Sung et al.	N/A	N/A
2006/0164268	12/2005	Lee et al.	N/A	N/A
2006/0256380	12/2005	Klassen et al.	N/A	N/A
2006/0277237	12/2005	Kuo et al.	N/A	N/A
2007/0041657	12/2006	Rychagov et al.	N/A	N/A
2007/0098283	12/2006	Kim et al.	N/A	N/A
2007/0206868	12/2006	Nakayama	N/A	N/A
2008/0068231	12/2007	Kuhns	N/A	N/A
2008/0253448	12/2007	Lin et al.	N/A	N/A
2009/0317008	12/2008	Cho et al.	N/A	N/A
2010/0260429	12/2009	Ichinose	N/A	N/A
2010/0303148	12/2009	Hiron et al.	N/A	N/A
2010/0332956	12/2009	Alrod et al.	N/A	N/A
2011/0019736	12/2010	Koyabu et al.	N/A	N/A
2011/0103703	12/2010	Karlov	N/A	N/A
2011/0150074	12/2010	Wang et al.	N/A	N/A
2012/0120287	12/2011	Funamoto et al.	N/A	N/A
2012/0224627	12/2011	Srinivasan	N/A	N/A
2012/0230595	12/2011	Kubo et al.	N/A	N/A
2012/0271802	12/2011	Oh	N/A	N/A
2012/0300792	12/2011	Patel et al.	N/A	N/A
2013/0002703	12/2012	Tripathi et al.	N/A	N/A
2013/0010864	12/2012	Teng	N/A	N/A
2013/0034309	12/2012	Nystad et al.	N/A	N/A
2013/0272626	12/2012	Robinson	N/A	N/A
2014/0002846	12/2013	Kondo	N/A	N/A
2014/0029846	12/2013	Su et al.	N/A	N/A
2014/0056577	12/2013	Ogawa et al.	N/A	N/A
2014/0092867	12/2013	Yu et al.	N/A	N/A
2014/0177730	12/2013	Chang	375/240.25	H04N 19/117
2014/0184612	12/2013	Dunaisky et al.	N/A	N/A
2014/0307962	12/2013	Seikh	N/A	N/A
2015/0103878	12/2014	Ramasubramonian et al.	N/A	N/A
2015/0138218	12/2014	Jeong et al.	N/A	N/A
2015/0138237	12/2014	Ghosh et al.	N/A	N/A

2015/0172670	12/2014	Li et al.	N/A	N/A
2015/0277776	12/2014	Okamoto et al.	N/A	N/A
2016/0006456	12/2015	Muramatsu et al.	N/A	N/A
2016/0057437	12/2015	Jeong et al.	N/A	N/A
2016/0179469	12/2015	Burgess et al.	N/A	N/A
2016/0321772	12/2015	Adsumilli	N/A	N/A
2016/0353117	12/2015	Seregin et al.	N/A	N/A
2017/0025098	12/2016	Keramidas et al.	N/A	N/A
2017/0177227	12/2016	Zhang et al.	N/A	N/A
2017/0230586	12/2016	Takahashi	N/A	N/A
2018/0054630	12/2017	Chang et al.	N/A	N/A
2018/0139446	12/2017	Abello et al.	N/A	N/A
2019/0027082	12/2018	Belle	N/A	N/A
2019/0325613	12/2018	Wu et al.	N/A	N/A
2020/0007149	12/2019	Fenney	N/A	N/A
2020/0007150	12/2019	Lacey et al.	N/A	N/A
2020/0007151	12/2019	Fenney et al.	N/A	N/A
2020/0007156	12/2019	Fenney	N/A	N/A
2020/0007866	12/2019	Martinelli et al.	N/A	N/A
2020/0090615	12/2019	Zhu et al.	N/A	N/A
2020/0252639	12/2019	Han et al.	N/A	N/A
2021/0006261	12/2020	Lacey et al.	N/A	N/A
2021/0036713	12/2020	Fenney et al.	N/A	N/A
2021/0067172	12/2020	Fenney	N/A	N/A
2022/0158653	12/2021	Lacey et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101106709	12/2007	CN	N/A
101283605	12/2007	CN	N/A
101653005	12/2009	CN	N/A
102036059	12/2010	CN	N/A
102289829	12/2010	CN	N/A
103210418	12/2012	CN	N/A
105379266	12/2015	CN	N/A
106095365	12/2015	CN	N/A
106993144	12/2016	CN	N/A
107567709	12/2017	CN	N/A
106233632	12/2018	CN	N/A
053010	12/1984	EP	N/A
0539010	12/1992	EP	N/A
1465418	12/2003	EP	N/A
2670139	12/2012	EP	N/A
3367683	12/2017	EP	N/A
2339989	12/1999	GB	N/A
2530312	12/2015	GB	N/A
2550965	12/2016	GB	N/A
9956239	12/1998	WO	N/A
1999056239	12/1998	WO	N/A
03096320	12/2002	WO	N/A
2003096320	12/2002	WO	N/A
2004091221	12/2003	WO	N/A
2008051350	12/2007	WO	N/A
2013006370	12/2012	WO	N/A
2014092867	12/2013	WO	N/A
2017199210	12/2016	WO	N/A
2018194815	12/2017	WO	N/A

OTHER PUBLICATIONS

Hal Canary: “Sk Image Info Reference”; Retrieved from the Internet:

URL:https://github.com/google/skia/blob/master/liteluserlapi/SkImageInfo_Reference.md; Jan. 26, 2018; 47 pages (note: copy in parent apn). cited by applicant

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS AND CLAIM OF PRIORITY

(1) This application is a continuation under 35 U.S.C. 120 of copending application Ser. No. 18/123,148 filed Mar. 17, 2023, now U.S. Pat. No. 12,064,934, which is a continuation of prior application Ser. No. 17/713,307 filed Apr. 5, 2022, now U.S. Pat. No. 11,611,754, which is a continuation of prior application Ser. No. 16/931,949 filed Jul. 17, 2020, now U.S. Pat. No. 11,323,718, which is a continuation of prior application Ser. No. 16/457,960 filed Jun. 29, 2019, now U.S. Pat. No. 10,757,415, which claims foreign priority under 35 U.S.C. 119 from United Kingdom Application No. 1810791.2 filed Jun. 29, 2018, the contents of which are incorporated by reference herein in their entirety.

BACKGROUND

(2) Data compression, either lossless or lossy, is desirable in many applications in which data is to be stored in, and/or read from, a memory. By compressing data before storage of the data in a memory, the amount of data transferred to the memory may be reduced. An example of data for which data compression is particularly useful is image data, such as depth data to be stored in a depth buffer, pixel data to be stored in a frame buffer and texture data to be stored in a texture buffer. These buffers may be any suitable type of memory, such as cache memory, separate memory subsystems, memory areas in a shared memory system or some combination thereof.

(3) A Graphics Processing Unit (GPU) may be used to process image data in order to determine pixel values of an image to be stored in a frame buffer for output to a display. GPUs usually have highly parallelised structures for processing large blocks of data in parallel. There is significant commercial pressure to make GPUs (especially those intended to be implemented on mobile devices) operate at lower power levels. Competing against this is the desire to use higher quality rendering algorithms on faster GPUs, which thereby puts pressure on a relatively limited resource: memory bandwidth. However, increasing the bandwidth of the memory subsystem might not be an attractive solution because moving data to and from, and even within, the GPU consumes a significant portion of the power budget of the GPU. The same issues may be relevant for other processing units, such as central processing units (CPUs), as well as GPUs.

(4) The embodiments described below are provided by way of example only and are not limiting of implementations which solve any or all of the disadvantages of known methods of data compression.

SUMMARY

(5) This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

(6) A method of converting 10-bit pixel data (e.g. 10:10:10:2 data) into 8-bit pixel data is described. The method involves converting the 10-bit values to 7-bits or 8-bits and generating error values for each of the converted values. Two of the 8-bit output channels comprise a combination of a converted 7-bit value and one of the bits from the fourth input channel. A third 8-bit output channel comprises the converted 8-bit value and the fourth 8-bit output channel comprises the error values. In various examples, the bits of the error values may be interleaved when they are packed into the fourth output channel.

(7) A first aspect provides a method of converting 10:10:10:2 format pixel data into 8888 format data, the method comprising, for each pixel: truncating the three 10-bit input data channels such that a first truncated channel and a third truncated channel each comprises 7 bits and a second truncated channel comprises 8 bits; appending a different bit from the 2-bit input data channel to each of the truncated first and third channels; generating an error value for each of the 10-bit input data channels by comparing, for each of the first and third channels, three LSBs of the input data channel and three bits derived from the truncated channel, and for the second channel, two LSBs of the input data channel and two bits derived from the truncated channel; outputting a first 8-bit channel comprising the truncated first channel with the appended bit from the 2-bit input data channel, a second 8-bit channel comprising the truncated second channel, a third 8-bit channel comprising the truncated third channel with the appended bit from the 2-bit input data channel and a fourth 8-bit channel comprising the three error values.

(8) A second aspect provides method of decompressing a block of compressed RGBA10:10:10:2 format pixel data, the method comprising: decompressing a compressed data block using a lossless decompression method, to generate a first decompressed data block, wherein the first decompressed data block comprises four channels of data and wherein the first, second and third channels are 8-bit channels; determining whether the fourth channel is an 8-bit channel and in response to determining that the fourth channel comprises fewer than 8-bits per pixel, converting the fourth channel into an 8-bit channel by adding one or more zeros to data for each pixel; removing an MSB from each of the first and third channels and combining them to form a 2-bit alpha channel; using bit replication to increase each of the first, second and third channels to 10-bit channels; for each of the first and third channels, modifying three LSBs by combining them with corresponding bits from the fourth channel to generate a 10-bit red output channel and a 10-bit blue output channel respectively and for the second channel, modifying two LSBs by combining them with corresponding bits from the fourth channel to generate a 10-bit green output channel.

(9) A third aspect provides a data compression unit arranged to convert 10:10:10:2 format pixel data into 8888 format data, the hardware comprising: Truncation logic arranged, for each pixel, to truncate the three 10-bit input data channels such that a first truncated channel and a third truncated channel each comprises 7 bits and a second truncated channel comprises 8 bits; Hardware logic arranged, for each pixel, to append a different bit from the 2-bit input data channel to each of the truncated first and third channels; Hardware logic arranged, for each pixel, to generate an error value for each of the 10-bit input data channels by comparing, for each of the first and third channels, three LSBs of the input data channel and three bits derived from the truncated channel, and for the second channel, two LSBs of the input data channel and two bits derived from the truncated channel;

Hardware logic arranged, for each pixel, to output a first 8-bit channel comprising the truncated first channel with the appended bit from the 2-bit input data channel, a second 8-bit channel comprising the truncated second channel, a third 8-bit channel comprising the truncated third channel with the appended bit from the 2-bit input data channel and a fourth 8-bit channel comprising the three error values.

(10) The data compression and/or decompression unit as described herein may be embodied in hardware on an integrated circuit. There may be provided a method of manufacturing, at an integrated circuit manufacturing system, a data compression and/or decompression unit as described herein. There may be provided an integrated circuit definition dataset that, when processed in an integrated circuit manufacturing system, configures the system to manufacture a data compression and/or decompression unit as described herein. There may be provided a non-transitory computer readable storage medium having stored thereon a computer readable description of an integrated circuit that, when processed, causes a layout processing system to generate a circuit layout description used in an integrated circuit manufacturing system to manufacture a data compression and/or decompression unit as described herein.

(11) There may be provided an integrated circuit manufacturing system comprising: a non-transitory computer readable storage medium having stored thereon a computer readable integrated circuit description that describes the data compression and/or decompression unit as described herein; a layout processing system configured to process the integrated circuit description so as to generate a circuit layout description of an integrated circuit embodying the data compression and/or decompression unit as described herein; and an integrated circuit generation system configured to manufacture the data compression and/or decompression unit as described herein according to the circuit layout description.

(12) There may be provided computer program code for performing any of the methods described herein. There may be provided non-transitory computer readable storage medium having stored thereon computer readable instructions that, when executed at a computer system, cause the computer system to perform any of the methods described herein.

(13) The above features may be combined as appropriate, as would be apparent to a skilled person, and may be combined with any of the aspects of the examples described herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Examples will now be described in detail with reference to the accompanying drawings in which:

(2) FIG. 1 shows a graphics rendering system;

(3) FIGS. 2A, 2B and 2C show three different data compression architectures;

(4) FIG. 3A is a flow diagram of an example lossy data compression method;

(5) FIG. 3B is a flow diagram of a data decompression method that may be used to decompress data that was compressed using the method of FIG. 3A;

(6) FIGS. 4A and 4B are schematic diagram showing different blocks of data and their subdivision into sub-blocks;

(7) FIG. 4C is a schematic diagram showing an example compressed data block;

(8) FIGS. 5A and 5B show two different example implementations of the analysis stage of the method of FIG. 3A;

(9) FIG. 6A is a flow diagram of a first example method of compressing a sub-block using the constant alpha mode of FIG. 3A;

(10) FIGS. 6B, 6C and 6D are schematic diagrams showing two different ways of packing the compressed values into a data block;

(11) FIG. 7A is a flow diagram of a second example method of compressing a sub-block using the constant alpha mode of FIG. 3A;

(12) FIG. 7B is a schematic diagram showing an example of how the pixels in a sub-block may be divided into the two subsets in the method of FIG. 7A;

(13) FIG. 8A is a flow diagram of a first example method of compressing a sub-block using the variable alpha mode of FIG. 3A;

(14) FIG. 8B is a schematic diagram illustrating a part of the method of FIG. 8A;

(15) FIG. 9 is a flow diagram of a second example method of compressing a sub-block using the variable alpha mode of FIG. 3A;

(16) FIG. 10A is a schematic diagram showing encoding patterns that may be used in the method of FIG. 9;

(17) FIGS. 10B and 10C are schematic diagrams showing two different ways in which compressed data for a mini-block is packed into a data field;

(18) FIG. 10D is a schematic diagram showing two further encoding patterns that may be used in the method of FIG. 9;

(19) FIG. 11 is a flow diagram of a further example method of compressing a sub-block using a constant alpha mode;

(20) FIG. 12A is a flow diagram of a method of converting an n-bit number to an m-bit number, where $n > m$;

(21) FIG. 12B is a schematic diagram of a hardware implementation of the method of FIG. 12A;

(22) FIG. 13A is a flow diagram of a method of converting an n-bit number to an m-bit number, where $n < m$;

(23) FIG. 13B is a schematic diagram of a hardware implementation of the method of FIG. 13A;

(24) FIGS. 13C and 13D are schematic diagrams illustrating two examples of the method of FIG. 13A;

(25) FIG. 14 is a flow diagram of a first example method of converting 10-bit data to 8-bit data;

(26) FIGS. 15A, 15B, 15C and 15D are schematic diagrams showing four different ways in which data may be packed into data fields dependent upon whether the methods of FIG. 9 or 11 are used;

(27) FIG. 16 is a flow diagram of a first example method of converting 10-bit data to 8-bit data;

(28) FIG. 17 is a schematic diagram that illustrates the method of FIG. 16;

(29) FIG. 18A is a flow diagram of a data compression method which combines the pre-processing method of FIG. 16 with a lossless data compression method;

(30) FIG. 18B is a flow diagram of a data decompression method that maybe used where data has been compressed using the method of FIG. 18A;

- (31) FIG. 19A is a schematic diagram of a further example data compression unit;
- (32) FIG. 19B is a flow diagram of a method of operation of the bit predictor element in the data compression unit shown in FIG. 19A;
- (33) FIG. 19C is a schematic diagram showing an example way in which data may be packed into a compressed data block where the method of FIG. 19A is used;
- (34) FIG. 20A shows a schematic diagram of another data compression unit;
- (35) FIG. 20B is a flow diagram of a method of lossless data compression;
- (36) FIG. 21A is a schematic diagram of a block of data that has been spatially decorrelated and remapped using the hardware of FIG. 20A;
- (37) FIG. 21B is a schematic diagram showing an encoded data output from the method of FIG. 22;
- (38) FIG. 22 is a flow diagram showing a method of entropy encoding;
- (39) FIG. 23 shows a computer system in which a data compression and/or decompression unit is implemented; and
- (40) FIG. 24 shows an integrated circuit manufacturing system for generating an integrated circuit embodying a data compression and/or decompression unit as described herein.
- (41) The accompanying drawings illustrate various examples. The skilled person will appreciate that the illustrated element boundaries (e.g., boxes, groups of boxes, or other shapes) in the drawings represent one example of the boundaries. It may be that in some examples, one element may be designed as multiple elements or that multiple elements may be designed as one element. Common reference numerals are used throughout the figures, where appropriate, to indicate similar features.

DETAILED DESCRIPTION

(42) The following description is presented by way of example to enable a person skilled in the art to make and use the invention. The present invention is not limited to the embodiments described herein and various modifications to the disclosed embodiments will be apparent to those skilled in the art.

(43) Embodiments will now be described by way of example only.

(44) As described above, memory bandwidth is a relatively limited resource within a processing unit (e.g. a CPU or GPU), similarly, memory space is a limited resource because increasing it has implications in terms of both physical size of a device and power consumption. Through the use of data compression before storage of data in a memory, both the memory bandwidth and the space in memory are reduced.

(45) Many data compression schemes exist, some of which are lossless and others that are lossy. Lossless compression techniques may be preferred in some situations because the original data can be perfectly reconstructed from the compressed data. In contrast, where lossy compression techniques are used, data cannot be perfectly reconstructed from the compressed data and instead the decompressed data is only an approximation of the original data. The accuracy of the decompressed (and hence reconstructed) data will depend upon the significance of the data that is discarded during the compression process. Additionally, repeatedly compressing and decompressing data using lossy compression techniques results in a progressive reduction in quality, unlike where lossless compression techniques are used. Lossless compression techniques are often used for audio and image data and examples of general purpose lossless compression techniques include run-length encoding (RLE) and Huffman coding.

(46) The amount of compression that can be achieved using lossless compression techniques (e.g. as described in UK patent number 2530312) depends on the nature of the data that is being compressed, with some data being more easily compressed than other data. The amount of compression that is achieved by a compression technique (whether lossless or lossy) may be expressed in terms of a percentage that is referred to herein as the compression ratio and is given by:

$$(47) \text{Compressionratio} = \frac{\text{Compressedsize}}{\text{Uncompressedsize}} \times 100$$

It will be appreciated that there are other ways to define the compression ratio; however, the above convention is used throughout. This means that a compression ratio of 100% indicates that no compression has been achieved, a compression ratio of 50% indicates that the data has been compressed to half of its original, uncompressed size and a compression ratio of 25% indicates that the data has been compressed to a quarter of its original, uncompressed size. Lossy compression techniques can typically compress data to a greater extent (i.e. achieve smaller compression ratios) than lossless compression techniques. Therefore, in some examples, e.g. where the extent of achievable compression is considered more important than the quality of the decompressed (i.e. reconstructed) data, lossy compression techniques may be preferred over lossless compression techniques. The choice between a lossless and a lossy compression technique is an implementation choice.

(48) The variability in the amount of compression that can be achieved (which is dependent upon characteristics of the actual data that is being compressed) has an impact on both memory bandwidth and memory space and may mean that the full benefit of the compression achieved is not realised in relation to one or both of these two aspects, as described below.

(49) In many use cases, random access of the original data is required. Typically for image data, to achieve this, the image data is divided into independent, non-overlapping, rectangular blocks prior to compression. If the size of each compressed block varies because of the nature of the data in the block (e.g. a block which is all the same colour may be compressed much more than a block which contains a lot of detail) such that in some cases a block may not be compressed at all, then in order to maintain the ability to randomly access the compressed data blocks, the memory space may be allocated as if the data was not compressed at all. Alternatively, it is necessary to maintain an index, with an entry per block that identifies where the compressed data for that block resides in memory. This requires memory space to store the index (which is potentially relatively large) and the memory accesses (to perform the look-up in the index) adds latency to the system. For example, in systems where it is important to be able to randomly access each compressed block of data and where an index is not used, even if an average compression ratio (across all data blocks) of 50% is achieved, memory space still has to be allocated assuming a 100% compression ratio, because for some blocks it may not be possible to achieve any compression using lossless compression techniques.

(50) Furthermore, as the transfer of data to memory occurs in fixed size bursts (e.g. in bursts of 64 bytes), for any given block there are only a discrete set of effective compression ratios for the data transfer to memory. For example, if a block of data comprises 256 bytes and the transfer of data occurs in 64 byte bursts, the effective compression ratios for the data transfer are

25% (if the block is compressed from 256 bytes to no more than 64 bytes and hence requires only a single burst), 50% (if the block is compressed into 65-128 bytes and hence requires two bursts), 75% (if the block is compressed into 129-192 bytes and hence requires three bursts) and 100% (if the block is not compressed at all or is compressed into 193 or more bytes and hence requires four bursts). This means that if a block of data comprising 256 bytes is compressed into anywhere in the range of 129-192 bytes, then three bursts are required for the compressed block, compared to four for the uncompressed block, making the effective compression ratio for the memory transfer 75% whilst the actual data compression achieved could be much lower (e.g. as low as 50.4% if compressed into 129 bytes). Similarly, if the compression can only compress the block into 193 bytes, the memory transfer sees no benefit from the use of data compression, as four bursts are still required to transfer the compressed data block to memory. In other examples, blocks of data may comprise a different number of bytes, and bursts to memory may comprise a different number of bytes.

(51) Described herein are various methods of performing data compression. Some of the methods described herein provide a guarantee that a compression threshold, which may be defined in terms of a compression ratio (e.g. 50%), compressed block size (e.g. 128 bytes) or in any other way, is met. An effect of this guarantee is that a reduced amount of memory space can be allocated whilst still enabling random access to blocks of compressed data and there is also a guaranteed reduction in the memory bandwidth that is used to transfer the compressed data to and from memory. In other examples the compression ratio may be targeted (i.e. the method may be configured to achieve the ratio in the majority of cases) but there is no guarantee that it will be met.

(52) Also described herein are methods for converting 10-bit (e.g. 10:10:10:2) data to 8-bit (e.g. 8:8:8:3) data and methods for mapping from an n-bit number to an m-bit number. As described below, the methods for converting 10-bit (e.g. 10:10:10:2) data to 8-bit (e.g. 8:8:8:3 or 8888) data may be used as a pre-processing (or pre-encoding) step for the methods of performing data compression described herein or may be used independently (e.g. with another data compression method or with only a lossless compression method, such as that described below with reference to FIGS. 20 A-B, 21A-B and 22). By first converting the 10-bit (e.g. 10:10:10:2) data using one of the methods described herein, the 10-bit can then subsequently be compressed by methods that are arranged to operate on 8888 format data. The conversion method may be lossy with respect to three of the channels (e.g. the RGB data) and lossless for the fourth channel (e.g. the alpha data); however as this format is typically used for high dynamic range (HDR) data and the majority of pixels (e.g. 75%) will still be of low dynamic range (LDR), the conversion can be performed with only a small loss of accuracy. The method for mapping from an n-bit number to an m-bit number described herein may be used within the methods of performing data compression as described below or may be used independently. By using this mapping method, data of other formats can be subsequently compressed by methods that are arranged to operate on 8888 format data and/or it can be used to reduce the internal buffering (e.g. registers, etc.) by, for example, 6 bits per pixel (i.e. 19%) and this may, for example, be used in the initial reserve compression sub-unit 204A described below and shown in FIG. 2C.

(53) FIG. 1 shows a graphics rendering system 100 that may be implemented in an electronic device, such as a mobile device. The graphics rendering system 100 comprises a host CPU 102, a GPU 104 and a memory 106 (e.g. a graphics memory). The CPU 102 is arranged to communicate with the GPU 104. Data, which may be compressed data, can be transferred, in either direction, between the GPU 104 and the memory 106.

(54) The GPU 104 comprises a rendering unit 110, a compression/decompression unit 112, a memory interface 114 and a display interface 116. The system 100 is arranged such that data can pass, in either direction, between: (i) the CPU 102 and the rendering unit 110; (ii) the CPU 102 and the memory interface 114; (iii) the rendering unit 110 and the memory interface 114; (iv) the memory interface 114 and the memory 106; (v) the rendering unit 110 and the compression/decompression unit 112; (vi) the compression/decompression unit 112 and the memory interface 114; and (vii) the memory interface 114 and the display interface 116. The system 100 is further arranged such that data can pass from the compression/decompression unit 112 to the display interface 116. Images, which are rendered by the GPU 104, may be sent from the display interface 116 to a display for display thereon.

(55) In operation, the GPU 104 processes image data. For example, the rendering unit 110 may perform scan conversion of graphics primitives, such as triangles and lines, using known techniques such as depth-testing (e.g. for hidden surface removal) and texturing and/or shading. The rendering unit 110 may contain cache units to reduce memory traffic. Some data is read or written by the rendering unit 110, to the memory 106 via the memory interface unit 114 (which may include a cache) but for other data, such as data to be stored in a frame buffer, the data preferably goes from the rendering unit 110 to the memory interface 114 via the compression/decompression unit 112. The compression/decompression unit 112 reduces the amount of data that is to be transferred across the external memory bus to the memory 106 by compressing the data, as described in more detail below.

(56) The display interface 116 sends completed image data to the display. An uncompressed image may be accessed directly from the memory interface unit 114. Compressed data may be accessed via the compression/decompression unit 112 and sent as uncompressed data to the display 108. In alternative examples the compressed data could be sent directly to the display 108 and the display 108 could include logic for decompressing the compressed data in an equivalent manner to the decompression of the compression/decompression unit 112. Although shown as a single entity, the compression/decompression unit 112 may contain multiple parallel compression and/or decompression units for enhanced performance reasons.

(57) In various examples, the compression/decompression unit 112 may implement a compression method (or scheme) that guarantees that a compression threshold (which may be pre-defined and hence fixed or may be an input variable) is met. As detailed above, the compression threshold may, for example, be defined in terms of a compression ratio (e.g. 50% or 25%), compressed block size (e.g. 128 bytes) or in any other way. In order to provide this guarantee in relation to the amount of compression that is provided, and given that the exact nature of the data is not known in advance, a combination of lossless and lossy compression methods are used and three example architectures are shown in FIGS. 2A-C. In most if not all cases, a lossless compression technique (such as that described in UK patent number 2530312 or as described below with reference to FIGS. 20A-B, 21A-B and 22) is used to compress a block of data and then a test is performed to determine whether the compression threshold is met. In the event that the compression threshold is not met, a lossy compression technique (such as vector quantisation (VQ) techniques) or the method described below with reference to FIGS. 3A and 4-11 that provides the guaranteed compression

according to the compression threshold) is instead applied to the data block to achieve the compression threshold.

(58) In the method shown in FIG. 2A, the uncompressed source data, (e.g. a block of 256 bytes) is input to both a primary compression unit **202** (which may also be referred to as a lossless compression unit) and a reserve compression unit **204** (which may also be referred to as a lossy or fallback compression unit). The input data block is therefore independently and in parallel compressed using two different methods (a potentially lossless method in the primary compression unit **202** and a lossy method in reserve compression unit **204**). An example method of lossless compression that may be implemented by the primary compression unit **202** is described below with reference to FIGS. 20A-B, 21A-B and 22. The reserve compression unit **204** compresses the input data block in such a way so as to guarantee that the compression threshold is satisfied. The two versions of the compressed data block are then input to a test and selection unit **206**. This test and selection unit **206** determines whether the compressed data block generated by the primary compression unit **202** satisfies the compression threshold (e.g. if it is no larger than 128 bytes for a 256 byte input block and a 50% compression threshold). If the compressed data block generated by the primary compression unit **202** satisfies the compression threshold, then it is output, otherwise the compressed data block generated by the reserve compression unit **204** is output. In all cases the compressed data block that is output satisfies the compression threshold and by only using lossy compression (in the reserve compression unit **204**) for those blocks that cannot be suitably compressed using lossless techniques (in the primary compression unit **202**), the overall quality of the compressed data is improved (i.e. the amount of data that is lost due to the compression process is kept low whilst still satisfying the compression threshold).

(59) In the method shown in FIG. 2B, the uncompressed source data, (e.g. a block of 256 bytes) is initially input to only the primary compression unit **202** and the input of the source data to the reserve compression unit **204** is delayed (e.g. in delay unit **208**). The amount of delay may be arranged to be similar to the time taken to compress the source data block using the lossless compression technique (in the primary compression unit **202**) or a little longer than this to also include the time taken to assess the size of the compressed data block output by the primary compression unit **202** (in the test and decision unit **210**). The compressed data block output by the primary compression unit **202** is input to the test and decision unit **210** and if it satisfies the compression threshold it is output and no lossy compression is performed. If, however, the compressed data block output by the primary compression unit **202** does not satisfy the compression threshold (i.e. it is still too large), then the test and decision unit **210** discards this compressed block and triggers the lossy compression of the block by the reserve compression unit **204**. The compressed data block output by the reserve compression unit **204** is then output.

(60) In the method shown in FIG. 2C, the reserve compression unit **204** is divided into two sub-units: an initial reserve compression sub-unit **204A** and a final reserve compression sub-unit **204B**, with each sub-unit performing a part of the lossy compression method. For example, the initial reserve compression sub-unit **204A** may compress each byte from 8 bits to 5 bits (e.g. using truncation or the method described below with reference to FIGS. 12A-B) and any further compression that is required to satisfy the compression threshold may be performed by the final reserve compression sub-unit **204B**. In other examples, the reserve compression sub-unit **204B** may perform a pre-processing step, (e.g. as described below with reference to FIG. 14). In yet further examples, the lossy compression method may be split in different ways between the two reserve compression sub-units **204A**, **204B**.

(61) In the method shown in FIG. 2C, the uncompressed source data, (e.g. a block of 256 bytes) is input to both the primary compression unit **202** and the initial reserve compression sub-unit **204A**. The input data block is therefore independently and in parallel compressed using two different methods (a lossless method in the primary compression unit **202** and the first part of a lossy method in sub-unit **204A**). The compressed data block output by the primary compression unit **202** is input to the test and decision unit **210** and if it satisfies the compression threshold it is output, the partially compressed data block output by the initial reserve compression sub-unit **204A** is discarded and no further lossy compression is performed for that data block. If, however, the compressed data block output by the primary compression unit **202** does not satisfy the compression threshold (i.e. it is still too large), then the test and decision unit **210** discards this compressed block and triggers the completion of the lossy compression of the block output by the initial reserve compression sub-unit **204A** by the final reserve compression sub-unit **204B**. The compressed data block output by the final reserve compression sub-unit **204B** is output.

(62) In certain situations, it may be possible to compress a data block by more than the compression threshold. In such instances, the primary compression unit **202** may output a compressed data block that always exactly satisfies the compression threshold or alternatively, the size of the output compressed data block may, in such situations, be smaller than that required to satisfy the compression threshold. Similarly, the lossy compression technique that is used in FIGS. 2A, 2B and 2C (and implemented in the reserve compression unit **204** or sub-units **204A**, **204B**) may output a compressed data block which always exactly satisfies the compression threshold or alternatively, the size of the compressed data block may vary whilst still always satisfying the compression threshold. In the case where a compressed data block is smaller than is required to exactly satisfy the compression threshold, there may still be memory bandwidth and memory space inefficiencies caused by fixed burst sizes and pre-allocation requirements respectively; however, as the compression threshold is satisfied, there is always an improvement seen in relation to both memory bandwidth and memory space. In various examples, headers may be used to reduce the used memory bandwidth for some blocks even further (e.g. by including in the header information about how much data to read from memory or write to memory).

(63) Depending upon the particular implementation, any of the architectures of FIGS. 2A-C may be used. The arrangement shown in FIG. 2A provides a fixed throughput and fixed latency (which means that no buffering of data is needed and/or no bubbles are caused later in the system) but the power consumption may be increased (e.g. compared to just having a single compression unit performing either lossless or lossy compression). The arrangement shown in FIG. 2B may have a lower power consumption (on average) than the arrangement shown in FIG. 2A because the reserve compression unit **204** in FIG. 2B can be switched off when it is not needed; however the latency may vary and as a result buffers may be included in the system. Alternatively, an additional delay element **208** (shown with a dotted outline in FIG. 2B) may be added between the test and decision unit **210** and the output to delay the compressed data block output by the primary compression unit **202** (e.g. the amount of delay may be arranged to be comparable to the time taken to compress the source data block using the lossy compression technique in the reserve compression

unit **204**). The inclusion of this additional delay element **208** into the arrangement of FIG. 2B has the effect of making the latency of the arrangement fixed rather than variable. The arrangement shown in FIG. 2C may also have a lower power consumption (on average) than the arrangement shown in FIG. 2A because the final reserve compression sub-unit **204B** in FIG. 2C can be switched off when it is not needed; however in some circumstances data is discarded by the initial reserve compression sub-unit **204A** that would have been useful later and this may reduce the accuracy of the decompressed data (for example, where data is compressed initially from 8 bits to 6 bits and then from 6 bits to 4 bits, the decompression from 4 bits back to 8 bits, may introduce more errors than if the data was compressed directly from 8 bits to 4 bits).

(64) The methods described above with reference to FIGS. 2A-C may be used in combination with any compression threshold; however, in many examples the compression threshold will be 50% (although this may be expressed in another way, such as 128 bytes for 256-byte data blocks). In examples where a compression threshold other than 50% is used, the compression threshold may be selected to align with the burst size (e.g. 25%, 50% or 75% for the example described above) and the architectures shown in FIGS. 2A-C provide the greatest efficiencies when this threshold can be met using lossless compression for the majority of the data blocks (e.g. >50%) and lossy compression is only used for the remainder of the blocks.

(65) To identify which compression technique was used (e.g. lossless or lossy), data may be appended that indicates the type of compression used (e.g. in a header) or this may be incorporated into any existing header (or header table) that is used or each compressed block of data may include a number of bits, in addition to the compressed data, that indicates the type of compression used (e.g. as described below with reference to FIGS. 3-11).

(66) In any of the architectures of FIGS. 2A-C, there may be an additional pre-processing step (not shown in FIGS. 2A-C) that is a lossy pre-processing step and puts the source data into a suitable format for the primary compression unit **202** and/or reserve compression unit **204**, **204A**. This lossy pre-processing step may, for example, change the format of the data from 10-bit (e.g. RGBA 1010102) format into 8-bit (e.g. RGBA8883 or 8888 format) and two example methods for performing this pre-processing are described below with reference to FIGS. 14 and 16. In various examples, the method of FIG. 16 may be used as a pre-processing step for the primary compression unit **202** and the method of FIG. 14 may be used as a pre-processing step for the reserve compression unit **204**, or vice versa, or the same method may be used for both the primary and reserve compression units.

(67) The use of different data formats and/or pre-processing steps in the architectures of FIGS. 2A-C may also require modifications to the compression methods used (e.g. in the primary compression unit **202** and/or reserve compression units **204**, **204A**, **204B**) and some examples of these are also described below. By combining a lossy pre-processing step with the lossless compression (implemented in the primary compression unit **202**), it will be appreciated that the compressed data which is output by the primary compression unit **202** is no longer lossless.

(68) A lossy compression technique which guarantees that a pre-defined compression threshold is met is described with reference to FIGS. 3-11. This technique may be implemented by the compression/decompression unit **112** shown in FIG. 1 and/or by the reserve compression unit **204** shown in FIGS. 2A and 2B. As described above, use of data compression reduces the requirements for both memory storage space and memory bandwidth, and guaranteeing that a compression threshold (which may be defined in any suitable way, as described above) is met ensures that benefits are achieved in terms of both memory storage space and memory bandwidth. In the examples described below the compression threshold is 50% (or 128 bytes where each uncompressed data block is 256 bytes in size); however in other examples the method may be used for different compression thresholds (e.g. 75% or 25%) and as described above the compression threshold selected may be chosen to correspond to an integer number of bursts for memory transfer.

(69) The lossy compression method shown in FIG. 3A takes as input, source data in RGBA8888 or RGBX8888 format or in corresponding formats with the channels in a different order (e.g. ARGB or other corresponding formats e.g. comprising four channels each having 8-bit values). The source data may, in various examples, comprise channels with data values having less than 8 bits and examples of the consequential changes to the method are described below (e.g. with reference to FIGS. 15A-D). In examples where the source data is not in a suitable format (e.g. where the RGB channels each comprise more than 8-bits), a pre-processing step (e.g. as described below with reference to FIG. 14 or FIG. 16) may be used to convert the source data into an appropriate format. Alternatively, the method of FIG. 3A may be used for data where the channels comprise more than 8-bits (e.g. 10:10:10:2 data); however by using the pre-processing technique described below with reference to FIG. 14 which includes an HDR flag, there is one extra bit that can be shared across the RGB values. The following examples relate to compressing and decompressing image data, e.g. in RGBA format, but it is to be understood that the same principles can be applied for compressing and decompressing other types of data in other formats.

(70) The source data that is input to the method of FIG. 3A comprises blocks of data. For image data each block of data relates to a tile (or block) of pixels (e.g. tiles comprising 8×8 pixels or 16×4 pixels) and each block is subdivided into a plurality of sub-blocks (block **302**). In various examples, each block of data is subdivided (in block **302**) into four sub-blocks. If the block of data is subdivided (in block **302**) into a smaller number of larger blocks, then the amount of compression that can be achieved may be larger but random access is made more difficult and unless many pixels in a block are accessed, the bandwidth usage increases as the 'data per accessed pixel' would increase. Similarly, with a larger number of smaller blocks, random access is made easier (and the data per accessed pixels may be reduced); however the amount of data compression that can be achieved may be reduced. If, for example, the block of data relates to an 8×8 tile of pixels or a 16×4 tile of pixels, the block may be subdivided into four sub-blocks **400** each corresponding to a 4×4 arrangement of pixels, as shown in FIG. 4A for an 8×8 tile and FIG. 4B for a 16×4 tile. The sub-blocks may be denoted sub-block 0-3. Having performed this sub-division (in block **302**), each sub-block is considered independently and a lossy compression mode is selected for each sub-block based on the results of an analysis of the alpha values for the pixels within the sub-block (block **304**). Dependent upon the outcome of this analysis, the selected mode may be a mode that uses a constant value for alpha (as applied in block **306** and referred to as the constant alpha mode) or a mode that uses a variable value for alpha across the sub-block (as applied in block **308** and referred to as the variable alpha mode). These may be the only two available modes or alternatively there may be one or more additional modes (e.g. as applied in block **310**). The compressed data for each sub-block (as output by one of blocks **306-310**) in a source data block is then packed together to form a

corresponding compressed data block (block 312).

(71) FIG. 4C shows a compressed data block 402 comprising compressed data 404 for each of the sub-blocks and a further data field 406 that indicates that the lossy compression method of FIG. 3A is being used. The data 404 for each sub-block 400 is divided into two fields: a 2-bit block mode 408 and a 252-bit block data 410. The block mode bits 408 indicate whether the variable alpha mode (block 308), constant alpha mode (block 306), or other mode (block 310) is used. The field values may, for example, be as follows:

(72) TABLE-US-00001 Field value Interpretation 0b00 Constant alpha 0b01 Variable alpha 0b1- Other modes (where used)

(73) An example implementation of the analysis stage (block 304) in FIG. 3A is shown in detail in FIG. 5A. In this example, the alpha values for each of the pixels within the sub-block are analysed and two parameters are computed: minalpha and maxalpha, which are the minimum and maximum values of alpha for all of the pixels in the sub-block (block 502). These may be determined in any way including, for example, use of a loop (as in the example pseudo-code below, or its functional equivalent) or use of a tree of tests, with the first step determining maximum and minimum alpha values for pairs of pixels and then the second step determining maximum and minimum alpha values for pairs of outputs from the first step, etc. These two parameters (minalpha and maxalpha) are then used in a subsequent decision process (blocks 504-508) and although the decision process is shown as being applied in a particular order, in other examples the same tests may be applied in a different order (e.g. blocks 506 and 508 may be swapped over, assuming alphadiff<254). Furthermore, it will be appreciated that the test in block 504 may alternatively be maxalpha>(minalpha+alphadiff).

(74) A first decision operation (block 504) assesses the range of alpha values across the sub-block and determines whether the range is greater than the errors that would be introduced by the use of the (best case) variable alpha mode (in block 308). The size of these errors is denoted alphadiff in FIG. 5A and this value may be predetermined. The value of alphadiff may be determined by comparing the loss in quality caused by the different methods within the variable alpha mode (i.e. 4-colour encoding with 4 bits of alpha or 3-colour encoding with 5 bits of alpha, and with two pixels sharing the same colour) in a training process (hence the use of the phrase 'best case' above). Alternatively, the value of alphadiff may be determined (again in a training process) by assessing different candidate values against a large test set of images to find the candidate value that provides the best results using either a visual comparison or an image difference metric. The value of alphadiff may be fixed or may be programmable.

(75) In response to determining that the range is greater than the errors that would be introduced by the use of the (best case) variable alpha mode ('Yes' in block 504), a variable alpha mode of compression (block 308) is applied to this sub-block. However, in response to determining that the range is not greater than the errors that would be introduced by the use of the (best case) variable alpha mode ('No' in block 504), a constant alpha mode of compression (block 306) is applied to this sub-block and two further decision operations (blocks 506, 508) are used to determine the value of alpha which is used for the entire sub-block. If the value of maxalpha is the maximum possible value for alpha (e.g. 0xFF, 'Yes' in block 506), then the value of alpha used in the constant alpha mode (constalphaval) is set to that maximum possible value (block 510). This ensures that if there are any fully opaque pixels, they stay fully opaque after the data has been compressed and subsequently decompressed. If the value of minalpha is zero (e.g. 0x00, 'Yes' in block 508), then the value of alpha used in the constant alpha mode (constalphaval) is set to zero (block 512). This ensures that if there are any fully transparent pixels, they stay fully transparent after the data has been compressed and subsequently decompressed. If neither of these conditions are held ('No' in both blocks 506 and 508), then an average value of alpha is calculated across the pixels in the sub-block (block 514) and used in the constant alpha mode.

(76) The following pseudo-code (or its functional equivalent) may, for example, be used to implement the analysis shown in FIG. 5 and in this code, P.alp is the alpha value for the pixel P being considered:

```
(77) TABLE-US-00002    CONST AlphaDiffTol = 4; U8 MinAlpha := 0xFF; U8 MaxAlpha := 0x00; U12 AlphaSum := 0;
FOREACH Pixel, P, in the 4x4block      MinAlpha := MIN(P.alp, MinAlpha);      MaxAlpha := MAX(P.alp, MaxAlpha);
      AlphaSum += P.alp; ENDFOR IF((MaxAlpha - MinAlpha) > AlphaDiffTol) THEN      Mode := VariableAlphaMode;
ELSEIF (MaxAlpha == 0xFF)      Mode := ConstAlphaMode;      ConstAlphaVal := 0xFF; ELSEIF (MinAlpha == 0x00)
      Mode := ConstAlphaMode;      ConstAlphaVal := 0x00; ELSE      Mode := ConstAlphaMode;
ConstAlphaVal := (AlphaSum + 8) >> 4; ENDIF
```

(78) An alternative example implementation of the analysis stage (block 304 in FIG. 3A) is shown in FIG. 5B. In this example, the parameter constalphaval is set initially to the alpha value of a pixel at a pre-defined location within the sub-block (block 520). For example, constalphaval may be set to the alpha value of the pixel at the top left of the sub-block (i.e. the first pixel in the sub-block). All the alpha values of the other pixels in the sub-block are then compared to this constalphaval (in blocks 522-526).

Where all the alpha values are very similar to constalphaval (e.g. within a range of ± 5 , 'Yes' in block 522) then the constant alpha mode (of block 306 in FIG. 3A) is used, but where they vary more than this ('No' in block 522) then the variable alpha mode (of block 308 in FIG. 3A) is used. Then, in a similar manner to the method of FIG. 5A, for the constant alpha mode, the parameter constalphaval is set to zero (in block 512) or the maximum value (in block 510) where the pixels are all nearly fully transparent (constalphaval<5, 'Yes' in block 524) or nearly fully opaque (constalphaval>250, 'Yes' in block 526) respectively. It will be appreciated that the particular values used in FIG. 5B as part of the analysis (e.g. in blocks 522-526) are provided by way of example only and in other examples these values may differ slightly.

(79) In comparison to the method of FIG. 5A, the method of FIG. 5B does not require the determination of minalpha and maxalpha which reduces the computational effort required to perform the analysis. However, the method of FIG. 5B may produce some visible artefacts (e.g. aliasing) particularly when an object moves slowly across the screen and is less likely to detect a 'constant alpha' tile because of the use of a pre-defined location as the centre of the alpha values.

(80) Where the analysis of the alpha values within a sub-block (in block 304, e.g. as shown in FIG. 5A or 5B) determines that the constant alpha mode (of block 306) is to be used and also sets the value of the parameter constalphaval (in one of blocks 510-514 and 520), then the compression of the sub-block proceeds as shown in FIG. 6A or FIG. 7A. FIG. 6A shows a flow diagram of a first example method of compressing a sub-block using the constant alpha mode (block 306 of FIG. 3A). For each pixel, each of

the RGB values are compressed from 8 bits to 5 bits (block **602** e.g. as described below with reference to FIGS. **12A-B** or using an alternative truncation approach) and then the compressed values are packed into a data block along with the value of constalphaval (block **604**). Therefore, in this example, the data for the 4×4 sub-block is compressed from 512 bits (in RGBA8888 format (16×32=512 bits)) to 248 bits (8+16×(5+5+5)=248 bits).

(81) Two different ways of packing the compressed values into a data block are shown in FIGS. **6B-6D**, although it will be appreciated that the compressed values may alternatively be packed into a data block in other ways. In the first example, as shown in FIG. **6B**, the data block **606** comprises three 80-bit fields **608-610**, each comprising data from one of the three channels (e.g. R, G or B) and a 12-bit field **612** which includes the 8-bit constalphaval (e.g. as determined using FIG. **5A** or **5B**) and the remaining 4 bits may be unused or reserved for future use. In each of the 80-bit fields **608-610** there are 5-bit values for each of the pixels in the sub-block.

(82) In the second example, as shown in FIGS. **6C-6D** the layout of data is similar to that for the second method of performing the variable alpha mode (of block **308**, as described below with reference to FIG. **9**) as this results in less complex hardware (and hence smaller hardware with a lower power consumption) where the two methods (i.e. the methods of FIGS. **6A** and **9**) are used together. As shown in FIG. **6C**, each sub-block **400** (e.g. each 4×4 pixel sub-block) is subdivided into four mini-blocks **650-653** (e.g. four 2×2 pixel mini-blocks). Each mini-block has a corresponding 60-bit data field **660-663** that contains the RGB data for the pixels in the mini-block. In FIG. **6C**, the mini-blocks and their corresponding data fields have been labelled Q, S, U and W. The 8-bit constalphaval (e.g. as determined using FIG. **5A** or **5B**) is distributed amongst the four 3-bit fields **670-673**. Within each of the mini-block data fields **660-663**, the RGB data is distributed as shown in FIG. **6D**. If each of the mini-blocks comprises four pixels, labelled A-D, these are each represented by three 5-bit values, one for each of the R, G and B channels (e.g. 5-bit values RA, GA and BA represent pixel A, 5-bit values RB, GB and BB represent pixel B, etc.).

(83) FIG. **7A** shows a flow diagram of a second example method of compressing a sub-block using the constant alpha mode (block **306** of FIG. **3A**) which is a variant of the method shown in FIG. **6A**. As shown in FIG. **7A**, the pixels in the sub-block **400** are divided into two non-overlapping subsets and then the pixels in each of the subsets are compressed by different amounts. In the specific example shown in FIG. **7A** which may be used where the constalphaval can be stored in less than 8 bits (e.g. the constant alpha is <5 or >250), the pixels in the first subset are compressed in the same way as in FIG. **6A**, i.e. by converting each of the R, G, B values from RGB888 format to RGB555 format (block **602**), whereas the pixels in the second subset are compressed in a different way (block **702**), i.e. by converting the RGB data from RGB888 format (i.e. three 8-bit values, one for each channel) to RGB565 format (i.e. 5-bit values for the R and B channels and 6-bit values for the G channel). The compression (in blocks **602** and **702**) may be performed as described below with reference to FIGS. **12A-B** or may use an alternative approach (e.g. truncation of the values by removing one or more of the LSBs). In other examples, the two subsets of pixels may be compressed in different ways (e.g. the pixels of the second subset may be compressed by converting the RGB data from RGB888 format to RGB554).

(84) FIG. **7B** shows an example of how the pixels in a sub-block may be divided into the two subsets. In the example shown in FIG. **7B**, the 4 pixels marked A-D form the first subset and the 12 shaded pixels form the second subset. In other examples, the split between the two subsets may be different (e.g. there may more or fewer than four pixels in the first subset, with the remaining pixels forming the second subset) and/or the position of the pixels in the first subset may be different.

(85) In examples where the constalphaval is an 8-bit value (e.g. where the constant alpha is not <5 or >250), the method of FIG. **7A** may be modified such that the pixels in the second subset are compressed in the same way as in FIG. **6A**, i.e. by converting each of the R, G, B values from RGB888 format to RGB555 format (block **602**), whereas the pixels in the first subset are compressed in a different way (block **702**), i.e. by converting the RGB data from RGB888 format (i.e. three 8-bit values, one for each channel) to RGB565 format (i.e. 5-bit values for the R and B channels and 6-bit values for the G channel). As before, the compression (in blocks **602** and **702**) may be performed as described below with reference to FIGS. **12A-B** or may use an alternative approach (e.g. truncation of the values by removing one or more of the LSBs).

(86) Where the analysis of the alpha values within a sub-block (in block **304**, e.g. as shown in FIG. **5A** or **5B**) determines that the variable alpha mode (of block **308**) is to be used, then the compression of the sub-block proceeds as shown in FIG. **8A** or **9**. FIG. **8A** shows a flow diagram of a first example method of compressing a sub-block using the variable alpha mode (block **308** of FIG. **3A**). As shown in FIG. **8A**, the data for each pixel in the sub-block is compressed by converting from RGBA8888 format, i.e. four 8-bit values, one for each channel including the alpha channel, to RGBA5555 format, i.e. four 5-bit values, one for each channel including the alpha channel, (block **802**, e.g. as described below with reference to FIGS. **12A-B**). The pixels in the sub-block **400** are then divided into two non-overlapping subsets (e.g. as described above with reference to FIG. **7B**) and the pixels in the first subset (e.g. pixels A-D in FIG. **7B**) are then subject to further compression (blocks **804-806**) which can be described with reference to FIG. **8B**. To further compress the pixels in the first subset, each of these pixels is compared to its neighbour pixels to identify which neighbouring pixel is most similar (block **806**). The similarity may be assessed using a colour difference parameter, with the most similar pixel having the smallest colour difference to the particular pixel and where colour difference between a pixel and a neighbour pixel (i.e. between a pair of pixels) may be calculated as:

(87)

$$.Math. Reddifference .Math. + .Math. Greendifference .Math. + .Math. Bluedifference .Math. + .Math. Alphadifference .Math. \quad (1)$$

(88) Having identified the most similar neighbouring pixel for each pixel in the first subset, an index is selected for each pixel in the first subset using a look-up table, such as:

(89) TABLE-US-00003 Index Most similar neighbouring pixel 000 R0 001 R1 010 R2 011 R3 100 R4

Where the references for the neighbouring pixels R0-R4 are defined as shown in FIG. **8B**. It will be appreciated however, that where the positions of the pixels of the first subset are different, these may be defined differently.

(90) In various examples, there may be additional indices that are used where there is a gradient around a pixel in the first subset. In such examples, the pixel data for one or more additional notional neighbour pixels are calculated, for example by the addition of a single additional notional neighbour pixel, R5, where the pixel data for R5 is calculated using:

(91) $R5 = (R0 + R1) / 2$

In other examples, one or more further notional neighbour pixels may also be considered, e.g. R6 and/or R7, where the pixel data for R6 and R7 is calculated using:

(92) $R6 = (R0 + R4) / 2$

$R7 = (R1 + R3) / 2$

(93) Where additional notional neighbour pixels are used, the look-up table includes corresponding additional entries to identify the indices, for example:

(94) TABLE-US-00004 Index Most similar neighbouring pixel 000 R0 001 R1 010 R2 011 R3 100 R4 101 R5 110 R6 111 R7

(95) In various examples, there may be an additional special case index that indicates that none of the neighbouring pixels (including any notional neighbouring pixels, where used) are sufficiently similar to the particular pixel. This may, for example, be determined based on a threshold and where the closest colour difference exceeds this threshold, an index of 000 may be used. In an example, the threshold may be 31 for 5-bit values. In addition to using the index 000, the pixel referred to by the index 000 is changed to a value that is an average between the current pixel and the pixel referred to. Alternatively, in a variation of FIG. 8A, if the conversion of pixel data from RGBA8888 to RGBA5555 (in block 802) is not performed until immediately prior to the packing of data values (in block 808), such that the comparison (in block 804) is performed on 8-bit values, the threshold will be different (e.g. 255 for 8-bit values).

(96) FIG. 9 shows a flow diagram of a second example method of compressing a sub-block using the variable alpha mode (block 308 of FIG. 3A). As shown in FIG. 9, the sub-block 400 (e.g. each 4x4 pixel sub-block) is subdivided (block 902) into four mini-blocks 650-653 (e.g. 2x2 pixel mini-blocks), as shown in FIG. 6C. Each mini-block is then compressed individually, with each mini-block having a corresponding data field 660-663 that contains the RGB data for the pixels in the mini-block. The corresponding four 3-bit fields 670-673 do not contain the constalphaval, as is the case in the earlier discussion of FIG. 6C, but where variable alpha mode is used, these 3-bit fields identify the encoding (or palette) mode that is used for each of the mini-blocks in the sub-block, as described below.

(97) The encoding mode is determined for each of the mini-blocks based on colour differences that are calculated for each pixel pair in the mini-block (block 904). The colour difference may be calculated using equation (1) above and this may be implemented by the functional equivalent of the pseudo-code provided below in which the colour difference is clamped to 6 bits (i.e. a maximum value of 63). In this code, the notation is as follows: IntermediateResult[5,0] refers to the 6 LSBs of the IntermediateResult 10-bit value and red/grn/blu/alp refer to red/green/blue/alpha respectively).

(98) TABLE-US-00005 U6 DiffMetric(PIXEL Pix1, PIXEL Pix2) { U6 Result; U10 IntermediateResult;
U8 R1 := Pix1.red; U8 R2 := Pix2.red; U8 G1 := Pix1.grn; U8 G2 := Pix2.grn; U8 B1 :=
Pix1.blu; U8 B2 := Pix2.blu; U8 A1 := Pix1.alp; U8 A2 := Pix2.alp; IntermediateResult:=
SAD4x5(R1, R2, G1, G2, B1, B2, A1, A2); IF((IntermediateResult > 63) THEN Result := 63;
ELSE Result := IntermediateResult[5..0]; ENDIF RETURN Result; }

(99) The pseudo-code above includes a sum of absolute differences (SAD) function and this may, for example, be implemented in any way (e.g. as implemented by a logic synthesis tool or as described in FIG. 2 of "Efficient Sum of Absolute Difference Computation on FPGAs" by Kumm et al).

(100) Having calculated the colour differences (in block 904), the smallest colour difference for any pixel pair in the mini-block is used to determine the mini-block encoding mode that is used (block 906). There are two distinct types of mini-block encoding mode that are used dependent upon whether the smallest colour difference (between any pixel pair in the mini-block) exceeds a threshold value (which may, for example be set at a value in the range 0-50, e.g. 40). If the smallest colour difference does not exceed the threshold ('Yes' in block 906), then one of a plurality of encoding patterns are used (as selected in block 908) and three per-mini-block palette colours are stored (blocks 910-914). However, if the smallest colour difference does exceed the threshold ('No' in block 906) then a four colour mode is used (block 916). These different mini-block modes are described in detail below. The encoding patterns rely on an assumption that in the majority of mini-blocks there are no more than three distinct colours and in such cases the mini-block can be represented by three palette colours along with an assignment of pixels to palette entries. The four colour mode is present to handle the exceptions to this.

(101) As noted above, the value of the threshold may be in the range 0-50. In various examples it may be a fixed value that is set at design time. Alternatively, it may be a variable which is stored in a global register and read each time the method of FIG. 9 is performed. This enables the threshold to be changed dynamically or at least periodically. In various examples, the value of the threshold may be set based on results from a training phase. The training phase may use an image quality metric, for example peak signal-to-noise ratio (PSNR) or structural similarity metric (SSIM) to assess a selection of images compressed using each of the three colour approach (blocks 908-912) or the four colour approach (block 916) and then the threshold value may be selected such that overall the highest image quality metrics are obtained.

(102) It may be noted that the threshold used in the method of FIG. 9 may have a different value to the threshold used in the method of FIG. 8A because the two thresholds serve different purposes as they are aiming to address different situations and remove different kinds of artefact. In the method of FIG. 8A, the threshold is used to identify when one of the pixels is an 'isolated colour', in that it cannot be represented well with one of its neighbours. In contrast, the threshold in the method of FIG. 9, the threshold is used to identify when the four colours are too different to each other.

(103) As shown in FIG. 9, if the smallest colour difference is smaller than or equal to the threshold ('Yes' in block 906), such that an encoding pattern can be used, the particular pattern that is used is selected from a set of six assignment patterns as shown in FIG. 10A based on the pixel pair (in the mini-block) that has the smallest colour difference (block 908). In each pattern in FIG. 10A, the two pixels that are shown shaded share a palette entry that is derived by averaging the two source pixel colours. Determining the three palette colours P.sub.0, P.sub.1, P.sub.2 (block 910) for the selected pattern (from block 908) therefore comprises performing this average and identifying the pixel data for the other two remaining pixels in the mini-block. Example pseudo-code for implementing this selection along with the calculation of three palette colours is provided below, but referring to

FIG. 10A, if the smallest colour difference is for pixel pair AB (denoted DiffAB in the pseudo-code), then the mode that is selected is the mode called 'Top', and the three palette colours are an average of A and B, and then pixels C and D. Similarly, if the smallest colour difference is for pixel pair CD (denoted DiffCD in the pseudo-code), then the mode that is selected is the mode called 'Bottom' and the three palette colours are an average of C and D and then pixels A and B. The palette colours in other modes (i.e. the 'Left', 'Right', 'Diag1' and 'Diag2' modes) are apparent from FIG. 10A.

(104) The following table shows an example mapping of the corresponding palette colours for each of the pixels A-D in a mini-block for each of the encoding modes along with the encoding value that is stored in the 3-bit field **670-673** for the mini-block (as shown in FIG. 6C). The pixels that are represented by identical palette colours are shown in bold and italic. In this example, the mapping has been arranged so that each pixel (e.g. pixel A, B, C or D) accesses only one of two possible palette colours and this results in a hardware implementation which is less complex than if all four pixels can access any of the three palette colours.

(105) TABLE-US-00006

Encoding Pixel A	Pixel B	Pixel C	Pixel D	value	Pattern name	(P.sub.0 or P.sub.1)	(P.sub.0 or P.sub.1)
(P.sub.0 or P.sub.2)	(P.sub.0 or P.sub.2)	000	Top	P.sub.1	P.sub.1	P.sub.0	P.sub.2
001	Bottom	P.sub.0	P.sub.1	P.sub.2	P.sub.2	010	
Left	P.sub.0	P.sub.1	P.sub.0	P.sub.2	011	Right	P.sub.1
P.sub.0	P.sub.1	P.sub.0	P.sub.2	P.sub.0	100	Diag 1	P.sub.0
P.sub.1	P.sub.0	P.sub.1	P.sub.2	P.sub.0	101	Diag 2	P.sub.1
P.sub.0	P.sub.0	P.sub.0	P.sub.2				

(106) A portion of example pseudo-code that implements this table (and blocks **908-910** of FIG. 9) is as follows:

(107) TABLE-US-00007

```

ELSIF (MinDiff == DiffAB)      MODE := TOP;      P0full:= PixelC;      P2full:=
PixelD;      P1full:= Average(PixelA, PixelB); ELSIF (MinDiff == DiffAC)      MODE := LEFT;      P0full:=
Average(PixelA, PixelC);      P1full:= PixelB;      P2full:= PixelD; ELSIF ...

```

(108) The Average function in the pseudo-code above may be implemented as follows (or its functional equivalent):

(109) TABLE-US-00008

```

PIXEL Average(PIXEL Pix1, PIXEL Pix2) {      PIXEL Result;      Result.red := (Pix1.red +
Pix2.red + 1) >> 1;      Result.grn := (Pix1.grn + Pix2.grn + 1) >> 1;      Result.blu := (Pix1.blu + Pix2.blu + 1) >> 1;
Result.alp := (Pix1.alp + Pix2.alp + 1) >> 1;      Return Result; }

```

(110) Having determined the three palette colours, P.sub.0, P.sub.1, P.sub.2 in RGBA8888 format, the values are converted to RGBA5555 format (block **912**, e.g. as described below with reference to FIGS. 12A-B). Whilst the operations of generating the palette colours (in block **910**) and the compression by converting format (in block **912**) are shown and described separately, it will be appreciated that the two operations may be combined into a single step and this may result in less complex hardware.

(111) As shown in FIG. 9, if the smallest colour difference is not smaller than or equal to the threshold ('No' in block **906**, i.e. the threshold is exceeded), then the encoding patterns shown in FIG. 10A cannot be used and instead the four colour mode is used. In this mode, the number of bits used to represent each channel is reduced (in block **916**) and in the example shown in FIG. 9, the pixel data is converted from RGBA8888 format to RGBA4434 format for each of the four pixels in the mini-block (block **916** e.g. as described below with reference to FIGS. 12A-B). The encoding value that is used to indicate that this mode is being used is one of the values that are not used for the encoding patterns (as detailed in the table above) and in various examples an encoding value of 111 may be used. This encoding value is stored in the 3-bit field **670-673** for the mini-block (as shown in FIG. 6C).

(112) Whilst the four colour mode is shown as using RGBA4434 format, alternative compressed formats may alternatively be used (e.g. RGBA4443). The RGBA4434 format may be used because although the human eye is sensitive to blue, the human visual system does not perceive the colour with the same resolution as red or green. In addition, a dither pattern may additionally be applied when using the four colour mode (e.g. in block **916** and also block **602** in FIG. 11).

(113) Having generated the compressed data for a mini-block, using any one of the encoding patterns or the four colour mode, the compressed data for the mini-block is packed into a 60-bit data field **660-663** for the mini-block and the corresponding encoding value is inserted in the associated field **670-673** (block **914**). The way that the data is packed differs depending upon the mode used and these are shown in FIGS. 10B and 10C. FIG. 10B shows the packing of pixel data for a mini-block that has been compressed using an encoding pattern and FIG. 10C shows the packing of pixel data for a mini-block that has been compressed using the four colour mode.

(114) As described above, for a mini-block that has been compressed using an encoding pattern, the data to include in the data field comprises three palette colours in RGBA5555 format and these may be packed into the data field as shown in FIG. 10B. As shown, the RGB data may be separated from the alpha channel data in order to more closely match the packing arrangement that may be used in the constant alpha mode (as shown in FIG. 6D). For a mini-block that has been compressed using the four colour mode, the data to include in the data field comprises four RGBA4434 values and these may be packed into the data field as shown in FIG. 10C. Again the arrangement shown is selected to more closely match the packing arrangement that may be used in the constant alpha mode. The arrangement for packing the data for each pixel is also shown in FIG. 10C and by packing the data in this way, with the four alpha bits being separated and placed between the RGB data, the R, G, B channel data corresponds to the positions of, respectively, four, four and three MSBs of the RGB data of the constant alpha mode shown in FIG. 6D. The alpha channel for the given pixel is distributed in the unused sections, matching LSB to MSB order.

(115) Having formed the data fields for the mini-blocks (in block **914**), these are then packed into data fields to the sub-blocks (block **918**, e.g. as shown in FIG. 6C).

(116) Although the encoding patterns are described above in relation to the variable alpha mode, in various examples they may also be used to provide a higher-precision constant alpha mode where the constant alpha value is 255 (or other pre-defined value) as shown in FIG. 11, with the previously described constant alpha mode (e.g. as shown in FIG. 6A or 7A) being used for other constant alpha values. FIG. 11 shows a flow diagram of a third example method of compressing a sub-block using the constant alpha mode (block **306** of FIG. 3A) which is a variant of the methods shown in FIGS. 6A and 9. In this example method, the sub-block **400** (e.g. each 4x4 pixel sub-block) is subdivided (block **902**) into four mini-blocks **650-653** (e.g. 2x2 pixel mini-blocks), as shown in FIG. 6C. Each mini-block is then compressed individually, with each mini-block having a corresponding data field **660-663** that contains the RGB data for the pixels in the mini-block. The corresponding four 3-bit fields **670-673** identify the encoding mode that is used for each of the mini-blocks in the sub-block. In this example, the constalphaval is not included as it is a known constant value, e.g. 255, and this in turn may be indicated using a dedicated block encoding mode (e.g. one of the 0b1-

modes which are indicated as reserved above). In yet further examples, the encoding patterns may in addition, or instead, be used to provide higher-precision constant alpha mode where the constant alpha value is stored and not assumed to have a known constant value (e.g. 255). In such examples, constalphaval may be stored alongside the pixel data and in order that there is room to store constalphaval, the method of FIG. 11 may be modified such that the three palette colours are converted from RGB888 to RGB676 (in block 1112, instead of RGB776) and the four-colour mode converts each colour from RGB888 to RGB554 (in block 602, instead of RGB555).

(117) The encoding mode is then determined for each of the mini-blocks based on colour differences that are calculated for each pixel pair in the mini-block (block 904). The colour difference may be calculated using any of the methods described above but omitting the alpha channel data.

(118) Having calculated the colour differences (in block 904), the smallest colour difference for any pixel pair in the mini-block is used to determine the mini-block encoding mode that is used (block 906). There are two distinct types of mini-block encoding mode that are used dependent upon whether the smallest colour difference (between any pixel pair in the mini-block) exceeds a threshold value (which may, for example be set at a value in the range 0-50, e.g. 40). If the smallest colour difference does not exceed the threshold ('Yes' in block 906), then one of a plurality of encoding patterns are used (as selected in block 908) and three per-mini-block palette colours are stored (blocks 910-914). However, if the smallest colour difference does exceed the threshold ('No' in block 906) then the earlier compression approach (from FIG. 6A), of converting each of the RGB values from 8 bits to 5 bits is used (block 602). As before the encoding patterns rely on an assumption that in the majority of mini-blocks there are no more than three distinct colours and in such cases the mini-block can be represented by three palette colours along with an assignment of pixels to palette entries. The earlier approach (of block 602) present to handle the exceptions to this. Where encoding patterns are used, higher precision can be used (e.g. more than 5 bits per colour channel) as there are only three palette colours rather than four pixels.

(119) The selection of an encoding pattern then proceeds as described above with reference to FIG. 9, except that there is no alpha data to consider (blocks 908-910). Having determined the three palette colours, P.sub.0, P.sub.1, P.sub.2 in RGB888 format, the values are converted to a format with higher resolution than RGB555, e.g. RGB776 or RGB666 format (block 1112, e.g. by truncation or as described below with reference to FIGS. 12A-B).

(120) Having generated the compressed data for a mini-block, using either one of the encoding patterns or the four colour mode (although in some examples, this four colour mode may not be used where alpha is constant value of 255), the compressed data for the mini-block is packed into a 60-bit data field 660-663 for the mini-block and the corresponding encoding value is inserted in the associated field 670-673 (block 1114). As described above, in various examples where the encoding patterns are used to provide higher-precision constant alpha mode, constalphaval may also be packed into the 60-bit data fields 660-663. Having formed the data fields for the mini-blocks (in block 1114), these are then packed into data fields to the sub-blocks (block 1118, e.g. as shown in FIG. 6C). This method of FIG. 11 provides improved RGB accuracy (e.g. RGB776 or RGB676 instead of RGB555) and hence improved image quality.

(121) In the description above, there are six different palette assignment (or encoding) modes, as shown in FIG. 10A (along with a four colour mode). In various examples there may be one or more additional mode, as shown in FIG. 10D. In such an additional mode there are only two palette colours (as indicated by the different shadings in FIG. 10D), with each of the two palette colour being the average of the two pixels sharing the same palette colour. In the first example, 1001, pixels A and C are averaged to generate the first palette colour, which is then used to replace the data for pixels A and C and similarly for pixels B and D to generate the second palette colour (which is then used to replace the data for those pixels). In the second example, 1002, pixels A and B are averaged to generate the first palette colour and pixels C and D are averaged to generate the second palette colour. As only two palette colours are stored (compared to the three in the earlier encoding modes), the palette colours may be stored at higher precision (e.g. RGBA8888 for constant alpha and RGBA7878 for variable alpha).

(122) The first of the two palette colour modes 1001 (as shown in FIG. 10D) may, for example be used if the colour difference between pixels A and C (DiffAC in the notation used above) and the colour difference between pixels B and D (DiffBD) are both less than a pre-defined threshold. Similarly, the second of the two palette colour modes 1002 may, for example be used if the colour difference between pixels A and B (DiffAB in the notation used above) and the colour difference between pixels C and D (DiffCD) are both less than a pre-defined threshold. Furthermore, whilst FIG. 10D shows only two two palette colour modes, it will be appreciated that there may, in various examples, be a third two palette colour modes in which one palette colour is an average of pixels A and D and the other palette colour is an average of pixels B and C.

(123) Where the two palette colour modes are used, an additional test against the pre-defined threshold (which may be different from the threshold used in block 906) may be inserted between blocks 904 and 906 in FIG. 9. If the test determines that a two palette colour mode may be used (as described above), the pixel data is averaged to form the two palette colours and then those palette colours may be converted to RGB7878 format (or they may be left as RGB888 for the constant alpha case) and then the pixel data is packed into a data field (block 914). If the test determines that a two palette colour mode cannot be used, then the method proceeds to block 906.

(124) As detailed above, there are many operations in data compression that involve converting an n-bit number, N, to an m-bit number, M, where $n > m$. The ideal mapping from n to m bits is given by:

$$(125) M = \text{IdealMap}(N, n, m) = \text{floor}\left(\frac{(2^m - 1)}{(2^n - 1)}N + \frac{1}{2}\right)$$

(126) This may, for example, be implemented using a large look-up table (LUT) comprising N entries. Alternatively, since the results are 'symmetrical' around a half, in that:

$$(127) \text{IdealMap}(\bar{N}, n, m) = \text{IdealMap}(N, n, m)$$

where X means the complement of the bits of X, then the size of the LUT may be halved. In such examples, one half of the LUT may be retained (e.g. the lower half) and the other half of the table (e.g. the upper half) may be derived by inverting 31 1 LSBs on input and similarly, inverting M-1 LSBs on output (e.g. by XORing X-1 LSBs with the MSB—the MSB is the same on input and output).

(128) FIG. 12A is a flow diagram of an alternative method of performing this conversion. Using the method of FIG. 12A provides a hardware implementation, as shown in FIG. 12B, that is both accurate and efficient (e.g. in hardware, i.e. in terms of size, etc.). Whilst it still uses a LUT, the LUT is significantly smaller than the LUTs mentioned above (but does require an additional addition element). Whilst the method is described herein in the context of data compression, the methods described herein with reference to FIG. 12A may also be used in other situations and independently of any other methods described herein.

(129) As shown in FIG. 12A, the method comprises receiving an input n-bit number, N, and truncating that number from n-bits to m-bits (block 1202) and this may be implemented in the truncation hardware unit 1222 in FIG. 12B. An adjustment value is then determined (block 1204) based on the input n-bit number, N, and, as described below, this may be implemented using a number of AND and OR gates in the look-up logic unit 1224. These AND and OR logic gates (or alternative logic arrangements that are functionally equivalent) compare a plurality of pre-determined subsets of the bits of the input n-bit number with pre-determined values in fixed-function circuitry and based on the outcome of the comparisons, determine an adjustment value which is then added to the truncated value (from block 1202 and truncation hardware unit 1222) in the increment/decrement unit 1226 (block 1206). The value of the adjustment value is either zero, one, or minus one.

(130) The operation of the look-up logic unit 1224 to determine the adjustment value (in block 1204) can be described with reference to a specific example where n=8 and m=5 and the following VHDL (or its functional equivalent), where “0-0-111” is an example of a pre-defined bit sequence which is compared to the bits in the input n-bit number and “01” its associated adjustment value:

```
(131) TABLE-US-00009    function CorrectionsFor8to5(i : std_logic_vector(7 downto 0))    return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin    results:=(others=>'0');    if std_match(i, "00---11-") then results:= "01";    end if;    if std_match(i, "0-0--111") then results:= results OR "01";    end if;    if std_match(i, "000--1-1") then results:= results OR "01";    end if;    if std_match(i, "0-00011-") then results:= results OR "01";    end if;    if std_match(i, "0--00111") then results:= results OR "01";    end if;    if std_match(i, "1--11000") then results:= results OR "11";    end if;    if std_match(i, "111--0-0") then results:= results OR "11";    end if;    if std_match(i, "1-11100-") then results:= results OR "11";    end if;    if std_match(i, "11---00-") then results:= results OR "11";    end if;    if std_match(i, "1-1--000") then results:= results OR "11";    end if;    return results; end function
```

CorrectionsFor8to5;

(132) In this example, ‘std_logic_vector(7 downto 0)’ refers to all 8 bits of N (labelled 0-7, with bit 0 being the LSB and bit 8 being the MSB) and the function ‘std_match’ may be implemented using AND gates, for example, ‘std_match(i, “00 - - - 11-”)’ may be implemented as follows:

b7Λb6Λb2Λb1

Where by is bit y of N and Λ represents a logical AND. Similarly ‘std_match(i, “11 - - - 00-”)’ may be implemented as:

b7Λb6Λb2Λb1

Furthermore, in this example, the initial value of ‘results’ is set to the two LSBs (i.e. bits 0 and 1) of N and the final value of ‘results’ is either set to 01 or 00 or to the initial value of results combined using an OR gate with a 2-bit value that is dependent upon the outcome of the ‘std_match’ function described above. This assumes that common factors are shared. It is estimated that approximately 31 AND/OR gates are required to implement the look-up logic unit 1224 for n=8 and m=5; however, it will be appreciated that any equivalent logic functionality (e.g. implemented using NANDs, NORs, XORs, etc.) may alternatively be used, e.g. dependent upon the silicon process used for fabrication or target clock speed).

(133) In this way, the look-up logic unit 1224 compares pre-determined subsets of the bits of the input n-bit number with pre-determined values in fixed-function circuitry, and sets the adjustment value in dependence on the results of the comparisons. Implementing the conversion from an n-bit number to an m-bit number in hardware using AND/OR gates as described above to determine an adjustment value to be applied to a truncated version of the n-bit number provides a very efficient implementation (e.g. in terms of silicon area, latency and power consumption) for performing a particular conversion (e.g. 8-bit values to 5-bit values). In this example, the hardware is inflexible in the sense that it is implemented in fixed-function circuitry using AND and OR gates for performing one or more specific conversions (e.g. conversions from 8-bit numbers to 5-bit numbers). If different conversions are required, e.g. different values of n and/or m, then different hardware can be used. The trade-off for this inflexibility is that the hardware can be implemented very efficiently (in terms of silicon area, latency and power consumption). Since the compression process described herein uses an 8-bit to 5-bit converter, and this is known at design time, the inflexibility is not a big problem whereas the efficiency advantages can be significant.

(134) Although data decompression is not described above, it will be appreciated that the data compression operations described above are reversed in order to achieve data decompression. For example, the methods of FIGS. 3A, 6A, 7A, 8A, 9 and 11 are reversed when performing data decompression. FIG. 3B shows a method of data decompression which is the reverse of the data compression method shown in FIG. 3A. As shown in FIG. 3B, for each sub-block of the compressed data, the compression mode is identified from the block mode (block 322) and depending upon the mode, the compressed sub-block data is decompressed using a variable alpha mode (block 308', which is the reverse of block 308 described above), or a constant alpha mode (block 306', which is the reverse of block 306 described above) and similarly if any other data compression mode is used (block 310'). The data decompression of the sub-blocks (in blocks 306', 308', 310') may involve converting a number from n-bits to m-bits where n<m (e.g. in inverse operations to blocks 602, 702, 802, 912, 916 and 1112). This may be implemented using the method shown in FIG. 13A and the hardware shown in FIG. 13B.

(135) As shown in FIG. 13A, the method comprises receiving an input n-bit number, N, and using bit replication to expand that number from n-bits to m-bits (block 1302) and this may be implemented in the replication hardware unit 1322 in FIG. 13B. This bit replication involves appending the (m-n) MSBs from N to the least significant end of N to generate an intermediate m-bit number, R, as shown in FIG. 13C. If m>2n, then (m-n)>n and hence it is necessary to first append all n-bits (i.e. the entirety of N) at least once (e.g. k times) and then append the (m-kn-n) MSBs from N to form the intermediate m-bit number R, as shown in FIG. 13D. The bit replication operation may alternatively be described as follows: assuming (nsm) and defining $k = \lfloor (m/n) \rfloor$ and $r = m \bmod n$, form a new value, R, comprising, in the MSB position, k repetitions of N, followed, in LSBs, the r MSBs of N.

(136) An adjustment value is then determined (block **1204**) based on the input n-bit number, N, and, as described below, this may be implemented using a number of AND and OR gates in the look-up logic unit **1224**. These AND and OR logic gates (or alternative logic arrangement that is functionally equivalent) compare a plurality of pre-determined subsets of the bits of N with pre-determined values in fixed-function circuitry and based on the outcome of the comparisons, determine an adjustment value which is then added to the replicated value (from block **1302** and replication hardware unit **1322**) in the increment/decrement unit **1226** (block **1306**). The value of the adjustment value is either zero, one, or minus one. In some examples, i.e. for some combinations of values of n and m (where $n < m$), the adjustment value is always zero and hence both the increment/decrement unit **1226** and the look-up logic unit **1224** can be omitted from FIG. **13B**.

(137) The operation of the look-up logic unit **1224** to determine the adjustment value (in block **1204**) can be described with reference to a specific example where $n=5$ and $m=8$ and the following VHDL (or its functional equivalent):

```
(138) TABLE-US-00010    function CorrectionsFor5to8(i : std_logic_vector(4 downto 0))    return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin    results:=(others=>'0');    if std_match(i, "00-11") then results:= "01"; end if;
    if std_match(i, "11-00") then results:= results OR "11"; end if;    return results; end function CorrectionsFor5to8;
```

(139) In this example, 'std_logic_vector(4 down to 0)' refers to all 5 bits of N (labelled 0-4, with bit **0** being the LSB and bit **4** being the MSB) and the function 'std_match' may be implemented using AND gates, for example, 'std_match(i, "00-11")' may be implemented as follows:

```
b4Ab3b1Ab0
```

Where by is bit y of N and A represents a logical AND. Similarly 'std_match(i, "11-00")' may be implemented as:

```
b4Ab3Ab1Ab0
```

Furthermore, in this example, the initial value of 'results' is set to the two LSBs (i.e. bits **0** and **1**) of N and the final value of 'results' is either set to 01 or 00 or to the initial value of results combined using an OR gate with 11. It is estimated that approximately 7 AND/OR gates are required to implement the look-up logic unit **1224** for $n=5$ and $m=8$.

(140) The look-up logic unit **1224** compares pre-determined subsets of the bits of the input n-bit number with pre-determined values in fixed-function circuitry, and sets the adjustment value in dependence on the results of the comparisons. As described above, the use of fixed-function hardware using the AND and OR gates described above allows the hardware to be implemented very efficiently in terms of silicon area, latency and power consumption. The hardware may be configured to perform conversions for a small group of different values for n and m.

(141) Further examples below show how the adjustment value is calculated (in block **1204** and unit **1224**) for other values of n and m.

(142) For $n=8$ and $m=4$, the following example VHDL shows how the adjustment value is calculated (in block **1204** and unit **1224**):

```
(143) TABLE-US-00011    function CorrectionsFor8to4(i : std_logic_vector(7 downto 0))    return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin    results:=(others=>'0');    if std_match(i, "00--11--") then results:= "01"; end if;
    if std_match(i, "000-1-1-") then results:= results OR "01"; end if;    if std_match(i, "0-0-111-") then results:= results OR "01"; end if;
    if std_match(i, "00001--1") then results:= results OR "01"; end if;    if std_match(i, "0-0011-1") then results:= results OR "01"; end if;
    if std_match(i, "00-01-11") then results:= results OR "01"; end if;    if std_match(i, "0--01111") then results:= results OR "01"; end if;
    if std_match(i, "1--10000") then results:= results OR "11"; end if;    if std_match(i, "11-10-00") then results:= results OR "11"; end if;
    if std_match(i, "1-1100-0") then results:= results OR "11"; end if;    if std_match(i, "11110--0") then results:= results OR "11"; end if;
    if std_match(i, "1-1-000-") then results:= results OR "11"; end if;    if std_match(i, "111-0-0-") then results:= results OR "11"; end if;
    if std_match(i, "11--00--") then results:= results OR "11"; end if;    return results; end function CorrectionsFor8to4;
```

It is estimated that approximately 41 AND/OR gates are required to implement the look-up logic unit **1224** for $n=8$ and $m=4$. In contrast, for $n=4$ and $m=8$ only bit replication is required and the adjustment value is zero in all cases.

(144) For $n=8$ and $m=3$, the following example VHDL shows how the adjustment value is calculated (in block **1204** and unit **1224**):

```
(145) TABLE-US-00012    function CorrectionsFor8to3(i : std_logic_vector(7 downto 0))    return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin    results:=(others=>'0');    if std_match(i, "00-11---") then results:= "01"; end if;
    if std_match(i, "0001-1--") then results:= results OR "01"; end if;    if std_match(i, "0-0111--") then results:= results OR "01"; end if;
    if std_match(i, "0001--11") then results:= results OR "01"; end if;    if std_match(i, "00-1-111") then results:= results OR "01"; end if;
    if std_match(i, "11-0-000") then results:= results OR "11"; end if;    if std_match(i, "1110-00") then results:= results OR "11"; end if;
    if std_match(i, "1-1000--") then results:= results OR "11"; end if;    if std_match(i, "11-00---") then results:= results OR "11"; end if;
    if std_match(i, "1110-0--") then results:= results OR "11"; end if;    return results; end function CorrectionsFor8to3;
```

It is estimated that approximately 31 AND/OR gates are required to implement the look-up logic unit **1224** for $n=8$ and $m=3$. In contrast, for $n=3$ and $m=8$ only bit replication is required and the adjustment value is zero in all cases.

(146) The compression methods described above with reference to FIGS. **2A-C**, **3A** and **4-11** take as input, source data in RGBA8888 or RGBX8888 format or in corresponding formats with the channels in a different order (e.g. ARGB). As mentioned above, in examples where the source data is not in this format, a pre-processing step may be used to convert the source data into an appropriate format (e.g. into 888Z where Z is an integer that is not greater than 8). A first example of this pre-processing step can be described with reference to FIG. **14** that converts data values from RGBA10:10:10:2 format to RGBA8883. An alternative technique for pre-processing 10-bit data is described below with reference to FIG. **16**. By using one or both of these pre-processing methods, the likelihood that the compression threshold is met (and hence a lossless compression method in the primary compression unit **202** of any of FIGS. **2A-C** is used) is increased; although minor modifications are required to the methods performed by the architectures shown in FIGS. **2A-C**, as described below. Furthermore, whilst these conversion methods are described herein in terms of a pre-processing step for the compression methods described above, both of the conversion methods

may alternatively be used independently of any of the other methods described herein. It will further be appreciated that whilst the compression method implemented by the primary compression unit **202** is itself lossless, if used in combination with the conversion method of FIG. **14** (which is not lossless), the overall data compression method is lossy. The conversion method of FIG. **16** is lossless and so is suited to be used as a pre-processing step for the primary compression unit **202** as the overall method is still lossless; however, if having performed the method of FIG. **16**, any bits from the fourth pseudo-channel are discarded, then the overall data compression method becomes lossy.

(147) FIG. **14** shows a flow diagram of a first example method of converting pixel data from RGBA 10:10:10:2 format to RGBA8883. As shown in FIG. **14**, the MSBs of each of the RGB channels are checked (block **1402**) and if one or more of these three MSBs are equal to one ('Yes' in block **1402**) then a flag is set (block **1404**), otherwise the flag is not set. This flag may be referred to as the high-dynamic-range (HDR) flag because if at least one MSB is equal to one then it is likely that the pixel data is HDR data. HDR images can represent a greater range of luminance levels than non-HDR images and HDR images are typically created by merging multiple low-or standard-dynamic-range (LDR or SDR) photographs or by using a special image sensor. Hybrid log-gamma is an HDR standard that defines a non-linear transfer function in which the lower half of the signal values (this is the SDR part of the range) use an x.sup.2 curve and the upper half of the signal values (this is the HDR part of the range) use a logarithmic curve and the reference white level is set to a signal value of 0.5. Of the 10 bits of R/G/B data, the most significant bit indicates whether a value is in the bottom half of the range (the SDR part) or in the top half of the range (the HDR part).

(148) As well as setting the flag or not, the pixel data is reduced from 10-bits to 8-bits in different ways dependent upon whether one or more of the MSBs for the RGB channels is one. If none of the three MSBs are equal to one ('No' in block **1402**), then each of the 10-bit values for the RGB channels is truncated by removing both the MSB (which is known to be a zero) and the LSB (block **1410**). If any of the three MSBs are equal to one ('Yes' in block **1402**), then there are two different ways in which the 10-bit values may be reduced to 8-bits (in block **1406**). In a first example, the two LSBs may be removed from each 10-bit value and in a second example, the method as described above with reference to FIGS. **12A-B** may be used.

(149) Where the method of FIG. **12A** is used, n=10 and m=8 and the following example VHDL shows how the adjustment value is calculated (in block **1204** and unit **1224**):

```
(150) TABLE-US-00013 function CorrectionsFor10to8(i : std_logic_vector(9 downto 0)) return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin results:=(others=>'0'); if std_match(i, "00-----11") then results:= "01"; end if; if
std_match(i, "0-00----11") then results:= results OR "01"; end if; if std_match(i, "0-0-00--11") then results:= results OR "01";
end if; if std_match(i, "0-0-0-0011") then results:= results OR "01"; end if; if std_match(i, "1-1-1-1100") then results:= results
OR "11"; end if; if std_match(i, "1-1-11--00") then results:= results OR "11"; end if; if std_match(i, "1-11----00") then results:=
results OR "11"; end if; if std_match(i, "11-----00") then results:= results OR "11"; end if; return results; end function
CorrectionsFor10to8;
```

It is estimated that approximately 25 AND/OR gates are required to implement the look-up logic unit **1224** for n=10 and m=8.

(151) To reverse this mapping (e.g. as shown in FIG. **13A**), such that n=8 and m=10, the following example VHDL shows how the adjustment value is calculated (in block **1204** and unit **1224**):

```
(152) TABLE-US-00014 function CorrectionsFor8to10(i : std_logic_vector(7 downto 0)) return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin results:=(others=>'0'); if std_match(i, "0011----") then results:= "01"; end
if; if std_match(i, "001-11--") then results:= results OR "01"; end if; if std_match(i, "001-1-11") then results:= results OR "01";
end if; if std_match(i, "110-0-00") then results:= results OR "11"; end if; if std_match(i, "110-00--") then results:= results OR
"11"; end if; if std_match(i, "1100----") then results:= results OR "11"; end if; return results; end function CorrectionsFor8to10;
```

It is estimated that approximately 19 AND/OR gates are required to implement the look-up logic unit **1224** for n=8 and m=10.

(153) Irrespective of the values of the three MSBs for the RGB channels for the pixel, the 2-bit alpha channel value is modified in the same way. As shown in FIG. **14**, the HDR flag is appended to the existing 2-bit value (block **1408**), making the output alpha channel value 3-bits.

(154) The method of FIG. **14** may be implemented on a per-pixel basis, although in variations on this method, the decision (in block **1402**) that leads to the setting of the HDR block may alternatively be performed less frequently, e.g. on a sub-block or mini-block basis.

(155) As noted above, where the method of FIG. **14** is used as a pre-processing step, the methods implemented in the architectures shown in FIGS. **2A-C** may be modified slightly. In particular, the compression threshold used by the test and selection unit **206** (in FIG. **2A**) or the test and decision unit **210** (in FIGS. **2B-C**) is calculated based on the size of the original data, prior to pre-processing, or alternatively, the threshold may be modified to account for the reduced size of the alpha channel values in the source data (i.e. 3-bits rather than 8-bits). This means, for example, that instead of defining the compression threshold in terms of a ratio, the threshold that is used may be defined in terms of the size of the compressed data block (e.g. 128 bytes) rather than a ratio. Additionally, where the reserve compression unit **204** implements the method described above with reference to FIGS. **3A** and **4-11**, this lossy compression method is also modified, as described below, as a consequence of the alpha value only comprising 3-bits, one of which is the HDR flag.

(156) In a first variation on the lossy compression method of FIG. **3A**, there may be three constant alpha modes that are used (in block **306**) along with the variable alpha mode (in block **308**). This is because there are only four possible alpha values, A, (instead of the 256 values where alpha is an 8-bit value) and hence the mode can be determined in a less complex manner than described above with reference to FIGS. **5A-B**. Each of the constant alpha modes has a different pre-defined value of constalphaval which can be specified by using a different field value that is inserted into the block mode bits **408** (in FIG. **4C**) and hence the value of constalphaval does not need to be stored separately (e.g. in fields **670-673** as is the case in the earlier description). As noted below, where the alpha value is zero, this uses the same field value as variable alpha; however as such a block is fully transparent, the RGB may be less important than for other alpha values and in this mode, the data is encoded as if alpha were variable per pixel.

(157) TABLE-US-00015 Field value Interpretation 0b00 Constant alpha Fully opaque: A = 0 × 3 0b01 Constant alpha

Transparent: A = 0 × 2 0b10 Constant alpha Transparent: A = 0 × 1 0b11 Variable alpha or Constant alpha, fully transparent: A = 0 × 0

(158) The constant alpha modes (of block **306**) are then implemented using the method of FIG. **11**, i.e. using the encoding patterns and three palette colours; however, in this case the palette colours are each represented as RGB676 plus a 1-bit HDR flag and hence block **1112** involves a different conversion to that described above. To perform this conversion, for the G channel, the value is truncated and for the R and B channels, the method described above with reference to FIGS. **12A-B** may again be used, with n=8 and m=6, or alternatively the values may simply be truncated.

(159) The adjustment value for n=8 and m=6 may be determined (in block **1204** and look-up logic unit **1224**) according to the following example VHDL (or its functional equivalent):

```
(160) TABLE-US-00016    function CorrectionsFor8to6(i : std_logic_vector(7 downto 0))    return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin    results:=(others=>'0');    if std_match(i, "00----11") then results:= "01";    end if;    if std_match(i, "0-00--11") then results:= results OR "01";    end if;    if std_match(i, "0-0-0011") then results:= results OR "01";    end if;    if std_match(i, "1-1-1100") then results:= results OR "11";    end if;    if std_match(i, "1-11--00") then results:= results OR "11";    end if;    if std_match(i, "11----00") then results:= results OR "11";    end if;    return results; end function CorrectionsFor8to6;
```

It is estimated that approximately 19 AND/OR gates are required to implement the look-up logic unit **1224** for n=8 and m=6.

(161) The reverse mapping (e.g. as shown in FIG. **13A**), such that n=6 and m=8, may be implemented according to the following example VHDL (or its functional equivalent):

```
(162) TABLE-US-00017    function CorrectionsFor6to8(i : std_logic_vector(5 downto 0))    return std_logic_vector is
variable results:std_logic_vector(1 downto 0); begin    results:=(others=>'0');    if std_match(i, "0011--") then results:= "01";    end if;    if std_match(i, "001-11") then results:= results OR "01";    end if;    if std_match(i, "110-00") then results:= results OR "11";    end if;    if std_match(i, "1100--") then results:= results OR "11";    end if;    return results; end function CorrectionsFor6to8;
```

(163) Furthermore, the four colour mode (in FIG. **11**) is modified such that the pixels are represented as RGB554 (instead of RGB555 as in block **602**, described above) along with a 1-bit HDR flag. The reduction of bit values from 8-bits to 5-bits or 4-bits may be performed as described above (e.g. with reference to FIGS. **12A-B**).

(164) Similar modifications in terms of data format are also made to the variable alpha mode, as described above with reference to FIG. **9**. In particular the three palette entries are encoded with a 1-bit HDR flag, 2 alpha bits and RGB665 (in blocks **910-912**) and the four colour mode uses a 1-bit HDR flag, 2 alpha bits and RGB444 (in block **916**).

(165) As a consequence of the different data formats used, the data is packed into the data fields (i.e. the 60-bit data fields **660-663** as shown in FIG. **6C**) in a slightly different way to that described above and this can be described with reference to FIGS. **15A-D**. FIGS. **15A** and **15B** relate to the constant alpha mode (as packed in block **1114** of FIG. **11**) and FIGS. **15C** and **15D** relate to the variable alpha mode (as packed in block **914** of FIG. **9**). FIGS. **15A** and **15C** show the packing where the four colour mode (of blocks **916** and **602**) are used and FIGS. **15B** and **15D** show the packing where the encoding patterns (of blocks **908-912** and **1112**) are used. The packing arrangement shown in FIG. **15C** replaces that shown in FIG. **10C** and described above. The packing arrangement shown in FIG. **15D** replaces that shown in FIG. **10B** and described above.

(166) FIG. **16** shows a flow diagram of a second example method of converting pixel data from 10-bit (e.g. RGBA 10:10:10:2) format to 8-bit format (e.g. 8888 format). This method may, for example, be used as a pre-processing step for RGBA 10:10:10:2 format data before it is input to the primary compression unit **202**. The method of FIG. **16** can be described with reference to the schematic diagram in FIG. **17**. As shown in FIG. **17**, the input pixel data comprises four input channels **1702-1708**, which are referred to here as RGBA channels although this data may alternatively represent other channels. The output from the method are four 8-bit pseudo-channels.

(167) The RGB input channels **1702-1706** each comprise 10 bits and the A (alpha) input channel **1708** comprises two bits. The RGB channels are each truncated by removing one or more LSBs (block **1602**) and in the method shown, the R and B channels are truncated to 7 bits **1710**, **1714** by removing three bits from each channel, **1711**, **1715** and the G channel is truncated to 8 bits **1712** by removing two bits **1713**. The 8-bit channel **1712** formed by truncation of the G channel forms one of four output pseudo-channels (block **1604**).

(168) The alpha channel data is optionally Gray coded (block **1606**) before one bit **1716** is appended as a new MSB on the truncated R data **1710** to form another one of the four output pseudo-channels **1720** and the other bit **1718** from the (optionally Gray coded) alpha channel data is appended as a new MSB on the truncated B data **1714** to form a further one of the four output pseudo-channels **1722** (block **1608**).

(169) The fourth of the four output pseudo-channels **1730** is formed from error values. In one example, to generate the error values, the truncated RGB787 data **1710**, **1712**, **1714** is replicated so that each channel comprises 10 bits **1724**, **1726**, **1728** (block **1610**). As shown in FIG. **17**, this bit replication involves copying the three (for R and B) or two (for G) MSBs and appending them as new LSBs to the truncated data. These new LSBs of the updated 10:10:10 data are then combined with a corresponding number of the original LSBs of the input data channels **1711**, **1713**, **1715** to generate the error values (block **1612**). One way to combine these three or two bit values would be using a subtraction operation (modulo 8 for R and B and modulo 4 for G), as follows:

```
(170) TABLE-US-00018 Red error = (3 new red LSBs - 3 LSBs of input red data) mod 8
Green error = (2 new green LSBs - 2 LSBs of input green data) mod 4
Blue error = (3 new blue LSBs - 3 LSBs of input blue data) mod 8
```

However, use of an XOR, as shown in FIG. **17**, provides a more compact hardware implementation. In other examples, any reversible function may be used. In an alternative example, instead of performing the bit replication of the MSBs as new LSBs to the truncated data, the three (for R and B) and two (for G) MSBs of the truncated RGB787 data may be directly combined with the original LSBs of the input data channels **1711**, **1713**, **1715** to generate the error values in block **1612**. As above, the combination may be through a subtraction, XOR, or any suitable reversible function.

(171) Having generated the three error values (in block **1612**), these are packed together to form the fourth pseudo-channel **1730** (block **1614**). This may be formed by concatenating the three error values, or alternatively, the bits from the error values may be

interleaved such that the three LSBs **1732** of the pseudo-channel **1730** comprise the LSB from each of the three error values, the two MSBs **1734** of the pseudo-channel **1730** comprise the MSB from each of the red and blue error values and the middle three bits **1736** comprise the remaining three bits, one from each of the three error values, as shown in FIG. **17**. By using this arrangement of bits in the fourth pseudo-channel **1730**, the reduction in image quality that results from the discarding of one or more of these bits (e.g. as described below with reference to FIG. **18A**) is minimised. For example, if the 3 LSBs of the fourth pseudo-channel **1730** are discarded, this results in the loss of a single MSB of the error value for each of the RGB channels, rather than the loss of the entire error value for one of the channels (e.g. the B channel).

(172) The following pseudo-code (or its functional equivalent) may be used to implement the method of FIG. **16**:

(173) TABLE-US-00019 $r[9:0] := \text{Replicate}(\text{Truncate}(r[9:0], 10, 7), 7, 10)$; $g[9:0] := \text{Replicate}(\text{Truncate}(g[9:0], 10, 8), 8, 10)$; $b[9:0] := \text{Replicate}(\text{Truncate}(b[9:0], 10, 7), 7, 10)$; $r_x[2:0] := r[2:0] \text{ xor } r[2:0]$; $g_x[1:0] := g[1:0] \text{ xor } g[1:0]$; $b_x[2:0] := b[2:0] \text{ xor } b[2:0]$; $a[1:0] := \{a[1], a[1] \text{ xor } a[0]\}$ $ch0_Red := \{a[0], r[9:3]\}$ $ch1_Grn := g[9:2]$ $ch2_Blu := \{a[1], b[9:3]\}$ $ch3_Lsb := \{r_x[2], b_x[2], r_x[1], g_x[1], b_x[1], r_x[0], g_x[0], b_x[0]\}$

(174) Whilst FIGS. **16** and **17** and the pseudo-code above show the bits from the alpha channels being appended as MSBs on the truncated R and B channel data (in block **1608**), in other examples the alpha bit may be inserted at different positions (e.g. as an LSB or in the middle of the truncated 7-bit data). However, the data compression that can be achieved, when this method is used in combination with a lossless compression technique, may be improved by positioning the alpha bit as the MSB.

(175) The four 8-bit pseudo-channels **1720**, **1712**, **1722**, **1730** generated using the method of FIG. **16** can then be fed into a data compression method that is designed to operate on RGBX8888 data, such as the lossless compression method implemented by the primary compression unit **202** in FIG. **2**.

(176) As described above, use of Gray coding of the alpha channel (block **1606**) is optional. By using Gray coding (in which only a single bit changes as the value increments by one), the amount of data compression that can be achieved using a lossless compression method (such as described below with reference to FIGS. **20A-B**, **21A-B** and **22** and/or in UK Patent Number GB2530312) is increased (e.g. by 1-2%). This is because it reduces the influence the alpha channel, which is often uncorrelated with the R, G or B data, has on both colour and/or spatial decorrelation steps, which are common first (or early) steps in lossless compression. Colour channel decorrelation often uses colour transforms such as:

(177) TABLE-US-00020 $R^* = R - \text{Gmod}2.\text{sup}.8$ $G^* = G$ $B^* = B - \text{Gmod}2.\text{sup}.8$ $A^* = A$

Where the pre-processing of FIG. **16** has been performed, the pseudo-channels do not correspond to the original colour data and instead the alpha data is combined into the red and blue channels. Where the value of alpha changes, the values of R^* and B^* can change significantly from one pixel to the next, making the data harder to compress; however, through the use of Gray coding, any single value change affects only R^* or B^* and not both.

(178) As was the case for the method of FIG. **14**, if the method of FIG. **16** is used as a pre-processing step, the methods implemented in the architectures shown in FIGS. **2A-C** may be modified slightly. In particular, where the method of FIG. **16** is used as a pre-processing step for the primary compression unit **202**, an additional option may be considered by the test and selection unit **206** in FIG. **2A** or the test and decision unit **210** in FIG. **2B** or **2C**. In the methods described above, these test and decision/selection units **206**, **210** compare the compressed data generated by the primary compression unit **202** to a compression threshold and if the compression threshold is satisfied, the compressed data generated by the primary compression unit **202** is output. If the threshold is not satisfied, data compressed by the reserve compression unit **204** (or sub-units **204A**, **204B**) is output instead. Where the pre-processing method of FIG. **16** is used, if the compression threshold is not satisfied by the compressed data generated by the primary compression unit **202**, the compressed data may be further reduced in size by discarding all of the data corresponding to the fourth pseudo-channel **1730** (i.e. the pseudo-channel formed from the error values). In such an implementation, the compressed data generated by the reserve compression unit **204** (or sub-units **204A**, **204B**) is only output if, having discarded the entirety of the data corresponding to the fourth pseudo-channel, the compression threshold is still not satisfied. This reduces the amount of data that is lost in data compression (since the lossy compression method in the reserve compression unit **204** is used less often) and hence improves the quality of the subsequently decompressed data.

(179) Instead of discarding the entirety of the fourth pseudo-channel, the bit prediction technique described below with reference to FIG. **19A**, may instead be used to partially discard bits prior to compression. In a further variation, a subset of the bits of the fourth pseudo-channel may be discarded and then the data may be fed back into the primary compression unit **202**.

(180) FIG. **18A** shows a flow diagram of a data compression method which combines the pre-processing method of FIG. **16** (block **1802**) with a lossless data compression method (block **1804**), such as described in GB2530312 and/or below with reference to FIGS. **20A-B**, **21A-B** and **22**. As noted above, some or all of the fourth pseudo-channel may be discarded (in block **1804**) to increase the amount of data compression. This discarding may be implemented, as described above, to ensure that a compression threshold is satisfied (as in FIGS. **2A-C**) or it may be implemented to improve compression but without any defined threshold or test and decision process. In order that the decompression hardware can operate in the same way irrespective of whether some or all of the fourth pseudo-channel data is discarded (in block **1804**), when discarding that data, one or more zeros (e.g. a single byte of data comprising all zeros) or another constant, pre-defined, sequence of bits, may be retained for that channel. Alternatively, a special packing code or other means may be used to indicate to the decompressor that the fourth pseudo-channel data has been discarded. The decompression can then proceed as described below with reference to FIG. **18B**.

(181) As shown in FIG. **18B**, the compressed data is initially decompressed (block **1810**) using the inverse of the lossless compression method used when compressing the data (in block **1802**). This operation outputs data in 888X format, as the fourth channel may be 8-bits or fewer, dependent upon how much data was discarded (as described above) and may be totally absent. The fourth channel is then reconstructed such that it comprises 8-bits for the pixel (block **1812**) and this may be implemented by adding zeros or another pre-defined sequence of bits. As described above, one or more zeros or the sequence of bits may be present in the compressed data and used in the reconstruction (in block **1812**) or alternatively, the reconstruction may proceed based on a known convention without any need for any bits relating to the fourth channel in the compressed data.

(182) The alpha channel data can then be extracted (block **1814**) by removing the MSBs from the first and third channels and

concatenating them (followed by optional Gray decoding, where this was used in the data compression), leaving 787 format data which is expanded to 10:10:10 format by bit replication (block **1816**). The error values from the fourth pseudo-channel are then applied using an inverse of the function used to generate those error values (in block **1612**). In the example shown in FIG. **18B**, the LSBs from each channel (e.g. 3 bits from the first and third channels and two bits from the second channel) are combined with the corresponding bits from the fourth channel using an XOR logic gate or equivalent logic block (block **1818**). Where any bits from the fourth channel were previously discarded, XORing with a zero will not change the value of the LSB. In examples where a pre-defined sequence of bits is used in the reconstruction that does not comprise all zeros, an alternative function may be used both to generate the error values when compressing the data (in block **1612**) and to apply the error values when decompressing the data (in block **1818**).

(183) In an example implementation, both the pre-processing methods of FIGS. **14** and **16** may be used. In such an example, the pre-processing method shown in FIG. **14** and described above may be used to pre-process RGBA10:10:10:2 format data before it is input to the reserve compression unit **204** or may be implemented prior to or by the initial reserve compression sub-unit **204A**. The pre-processing method shown in FIG. **16** and described above may be used to pre-process RGBA 10:10:10:2 format data before it is input to the primary compression unit **202**.

(184) The methods described above are described with reference to RGBA or RGBX data but as noted above, the methods may also be used with data other than image data, where three of the four channels (e.g. the RGB channels above) are correlated across channels in some way (e.g. colour or spatially correlated). For some types of data, such as YUV or YCbCr data (where YCbCr data is digital luma and two chroma components and YUV is its analogue counterpart, although it may also be used to mean YCbCr) which are frequently arranged in two or three planar formats (e.g. dependent upon whether the U and V data is separate or interleaved), there may be little or no correlation across channels within a data block, as the channels are often assembled by gathering disjoint, albeit adjacent, single-or dual-channel pixel blocks e.g. multiple tiles of only Y, only U or only V or only UV are packed as the channels of the block. As a consequence, whilst the compression methods described above can be used, and despite that the palette modes of the lossy compression method (e.g. as described above with reference to blocks **908-912** of FIG. **9**) work well where the data has strong spatial correlation as well as good correlation between channels, they do not work well for YUV or YCbCr data (or other uncorrelated data types) as it is difficult to predict a value from its neighbours and therefore it relies too heavily on the four colour mode (block **916** of FIG. **9**).

(185) FIG. **19A** is a schematic diagram of a data compression unit which increases the amount of compression that is achieved compared to a primary compression unit **202** that implements only a lossless compression method such as the method described in UK Patent Number GB2530312 or the method described below with reference to FIGS. **20A-B**, **21A-B** and **22**. This data compression unit shown in FIG. **19A** may be used as the primary compression unit **202** in the architectures of any of FIGS. **2A-C**, particularly where the input data is not correlated across channels and therefore cannot greatly benefit from colour decorrelation or colour palettes. Furthermore, as the data compression unit of FIG. **19A** may operate separately on each channel, it may be implemented within the primary compression unit **202** for only a subset of the channels (e.g. for the alpha channel or the fourth pseudo-channel created using the method of FIG. **16**) whilst the remaining channels are compressed using a lossless compression method such as the method described in UK Patent Number GB2530312 or the method described below with reference to FIGS. **20A-B**, **21A-B** and **22**. In examples where the data compression unit of FIG. **19** is used as the primary compression unit **202** in the architectures of any of FIGS. **2A-C**, there is less reliance upon the reserve compression unit **204** which may involve more severe quantisation and hence result in a larger reduction in image quality.

(186) The data compression unit shown in FIG. **19A** may alternatively be used separately from the architectures shown in FIGS. **2A-C**, i.e. without any test and selection/decision unit **206**, **210** or reserve compression unit **206** and where there is no need to guarantee that a compression threshold is met. Instead, by using the data compression unit shown in FIG. **19A**, the amount of data compression achieved overall may be increased and a particular compression ratio may be targeted. The data compression unit shown in FIG. **19A** may also be used with compression methods other than those described in UK Patent Number GB2530312 or the method described below with reference to FIGS. **20A-B**, **21A-B** and **22** and/or for compression of data other than image data (e.g. for any data that has a compression algorithm as well as multiple levels of quantisation, where further quantisation improves the efficiency of the compression algorithm).

(187) As shown in FIG. **19A**, source data block (which for the purposes of the following description is assumed to be in **8888** format) is split up and each channel may be considered separately (as indicated by the dotted box in FIG. **19A**). Alternatively, two or more channels (or any subset of the channels) may be considered together (and in this case, the total size that is determined relates to all channels being considered together). For example, RGB data may be considered together, with a single level of quantisation (i.e. a single mapping parameter) being chosen uniformly across the three colour channels. The following description refers to consideration of each channel separately by way of example only.

(188) The data for a single channel is input to both a bit predictor element **1902** and an x-bit remapping element **1904**. The operation of the x-bit remapping element is controlled based on an input from the bit predictor element **1902**, referred to as the mapping parameter. This mapping parameter specifies a data conversion or truncation operation that is performed by the x-bit remapping element **1904** and in particular, this mapping parameter specifies the value x. By selecting a lower value of x, the amount of data compression that is achieved is increased, however, the accuracy is reduced since the remapping is a lossy operation. The data conversion operation that is performed by the x-bit remapping element **1904** may also include operations such as low pass filtering where the mapping parameter, x, determines how heavily filtered the data is.

(189) The operation of the bit predictor element **1902** is shown in FIG. **19B**. The bit predictor element **1902** takes the data for the channel (e.g. 8-bits per pixel for each pixel in the data block), compresses it using the same data compression method used in the data compression element **1906** (block **1922**), however, some operations of the data compression method may be omitted (e.g. those operations after entropy encoding). For example, the bit predictor element **1902** may not pack the compressed data into a compressed data block but may instead just determine the total size of the compressed data (block **1924**). This total size is then compared to a plurality of thresholds (e.g. A.sub.Y, B.sub.Y, C.sub.Y, D.sub.Y) which may, for example, be stored in a look-up

table, and based on the comparisons, a mapping parameter (i.e. a level of quantisation and hence a value of x) is determined (block **1926**) and output.

(190) In examples where the data compression within the bit predictor element **1902** is performed using a lossless compression method such as the method described in UK Patent Number GB2530312 or the method described below with reference to FIGS. **20A-B**, **21A-B** and **22**, the bit predictor element **1902** may compress the data (in block **1922**) using wavelet transformation, sign remapping and group length generation; however, group length codes, packing mode codes and the actual packing of data may be omitted as they are not needed in order to generate the size of the compressed block. In such examples, the total size of the compressed data for the channel may be given by the size of the variable length data of the channel (see the description of FIGS. **20A-B**, **21A-B** and **22** below). This size may be calculated by summing four times the coefficient length for each group (which comprises a 2×2 set of values) other than the top-left. The group sizes may be determined by combining neighbouring values into groups of four and then the number of bits required to encode the largest of these is used as the group bit-depth and is encoded in the block. The size of the group is then four times this bit-depth. The top left group is treated differently as the reference value is always stored at full precision and its size may be determined from a look-up table. In other examples, any suitable data compression method may be used within the bit predictor element **1902** and the data compression element **1906**.

(191) The comparison to the plurality of thresholds (in block **1926**) may, for example, be implemented as set out in the following pseudo-code (or its functional equivalent), where Chan Size is the total size of the compressed data for the channel:

(192) TABLE-US-00021 If (Chan Size $\leq$ A.sub.Y) $x=8$ Else if (Chan Size $\leq$ B.sub.Y) $x=7$ Else if (Chan Size $\leq$ C.sub.Y) $x=6$ Else if (Chan Size $\leq$ D.sub.Y) $x=5$ Else $x=4$

and in various examples, a look-up table may be used. The pseudo-code (or look-up table) comprises at least two rows but may comprise many more rows. Furthermore, whilst the value of x is described herein as referring to an actual bit depth, in other examples the mapping parameter, x , may correspond to a set of different quantisations.

(193) The values of each of the plurality of thresholds (e.g. A.sub.Y, B.sub.Y, C.sub.Y, D.sub.Y) used by the bit predictor element **1902** may be determined by compressing a large training set of images multiple times using different configurations of thresholds and analysing the resultant sizes of the compressed data, or the score of an image quality metric for a given target compression ratio. The thresholds may be set to achieve a particular compression threshold or target compression ratio or may be calibrated to minimise certain artefacts of loss compression such as block artefacts or banding. The thresholds may also be set such that only those blocks that do not compress well are quantised.

(194) Having determined the mapping parameter (i.e. value of x) in the bit predictor element **1902**, the x -bit remapping element **1904** converts the input data from its original format, which may be 8-bits per pixel to x -bits per pixel, where x is the mapping parameter. This conversion may, for example, use truncation or may alternatively use the method described above with reference to FIGS. **12A-B**. As described above, in other examples, the mapping parameter may correspond to a set of different quantisations rather than an actual bit depth and in such examples, the remapping element **1904** converts the input data based on the mapping parameter. In other examples, the bit depth might not be altered but instead, the information content may be reduced by low pass filtering, under the control of the mapping parameter, x . Alternatively, a combination of both filtering and bit precision reduction may be employed.

(195) The data that is generated by the x -bit remapping element **1904** is input to the data compression element **1906** that then compresses the data using a data compression method, such as the lossless compression method described in UK Patent Number GB2530312 or the method described below with reference to FIGS. **20A-B**, **21A-B** and **22**. It will be appreciated that whilst this method may be described as lossless and may be lossless based on the data input to the data compression element **1906**, the resultant compressed data is not lossless with reference to the source data, since the conversion performed by the x -bit mapping element **1904** may be lossy.

(196) In order that the data can subsequently be decompressed, it is necessary to include the value of x the compressed data in some way, so that the correct inverse mapping may be used. FIG. **19C** shows an example of the way in which data for the four channels may be packed into a compressed data block (in block **1910**). As shown in FIG. **19C**, the size of the data for each channel **1930-1933** may be different where the data compression method includes entropy encoding. The 8-bit mode field **1940** comprises a 2-bit packing mode for each channel. One of the possible packing modes may be used to indicate that $x < 8$; however, given the available number of bits, it may not be possible to include the value of x for the channel in the packing mode without increasing the size of the mode field **1940**. Consequently, the bit depth encoding element **1908** may encode the value of x (where it is less than 8) in the data generated by the data compression element **1906**. For example, the value of x may be encoded within the reference value field of the compressed channel, e.g. by changing the least significant unused bit from 0 to 1, i.e. a unary encoding. As described below with reference to FIGS. **20A-B**, **21A-B** and **22**, the reference value is a data value which is not entropy encoded but is included in a specified reference value field of the encoded data (as generated by the data compression element **1906**). This reference value may, for example, be modified to encode the value of x as follows (where '0b' is a binary value indicator and e,f,g,h,i,j,k are used to identify bits that are used to encode the reference value):

(197) TABLE-US-00022 x Original value Encoding 7 0b0efghijk 0b1efghijk 6 0b00fghijk 0b01fghijk 5 0b000ghijk 0b001ghijk 4 0b0000hijk 0b0001hijk

(198) If the compressed data has no redundancy that can be exploited, then the value x may instead be packed alongside the compressed data, e.g. occupying a header byte. This can be read at the point of decompression to determine which inverse mapping to perform after decompressing the block in order to map the data back up to its original precision.

(199) FIG. **19A** also shows an optional y -bit remapping element **1912**. This may be used to reduce the number of bits per pixel in the source data for the channels that are used by the bit predictor element **1902**. The y -bit remapping element **1912** may use truncation or the method described above with reference to FIGS. **12A-B** to reduce the bit depth of the source data. The bit depth, i.e. the value of y , may be a value between 4 and 8, where $y=8$ provides the most accurate predictions (i.e. in terms of selecting a value of x to satisfy an overall compression threshold or target) but is most expensive in terms of operations performed by the bit predictor element **1912**. A value of $y=4$ provides the least accurate predictions but is least expensive in terms of hardware and

operations implemented by the bit predictor element **1902**. Similar to the value x , the value y may correspond to some other kind of quantisation, e.g. a low pass filter. By removing noise from the data the accuracy of the predictor's result can be improved, and can be less susceptible to change due to small perturbations in the original data.

(200) It will be appreciated that whilst the description of FIGS. **19A-C** above refer to source data comprising 8-bits per pixel per channel, the methods may also be used where there are more or fewer bits per pixel per channel, e.g. 10, and this may result in more or fewer thresholds being used (in block **1926**).

(201) In examples where the source data is 10-bit data, rather than 8888 data, the data compression unit may include an additional pre-processing element, not shown in FIG. **19A**. This pre-processing element may perform a data conversion similar to that described above with reference to FIGS. **16** and **17**. The 10-bit values are truncated and then expanded back to 10-bit values using bit replication, thereby generating new LSBs and error values are generated by XORing the LSBs of each original 10-bit value with the new LSBs of the corresponding truncated expanded value. The 8 MSBs of three tiles (which correspond to three input channels) form the first, second and third pseudo-channels whilst the error values are packed into the fourth pseudo-channel as described above and shown in FIG. **17**. One set of thresholds is then used in the bit predictor element **1902** for the first three pseudo-channels and these thresholds may be different from those used for 8888 data. A second set of thresholds (which also may be different from those used for 8888 data) is used for the fourth pseudo-channel and the quantisation options may be reduced, e.g. to 6 and 3 bits only. Where the data compression unit of FIG. **19A** is used as the primary data compression unit **202** in any of the architectures of FIGS. **2A-C**, the test and selection/decision units also operate differently. It is initially checked whether all four pseudo-channels fit within the target size, in which case they are included. If they do not all fit, it is checked whether the first three pseudo-channels fit with a single additional byte, and if so the fourth pseudo-channel is replaced by a constant zero channel (occupying a single byte). If this test still does not satisfy the target compression threshold, then the output from the reserve compression unit is used. The data can subsequently be decompressed using a modified version of the method described above with reference to FIG. **18B**.

(202) In examples where the data compression unit shown in FIG. **19A** targets a particular compression rate (e.g. 50%) and uses the lossless compression method described in UK Patent Number GB2530312 or the method described below with reference to FIGS. **20A-B**, **21A-B** and **22**, the RAW mode of the compression method may be modified to store data at a different precision, e.g. 4 bits rather than 8 bits.

(203) The data compression unit shown in FIG. **19A** is a versatile unit that is able to compress multiple types of data, e.g. YUV, RGB(A), non-image data, etc. It is applicable to various combinations of quantisation/lossy compression and can be used to improve compression rates or target a specific compression ratio. The thresholds that are used may be calibrated to produce desired properties, such as minimising block artefacts or banding or to optimise a given metric score, and the thresholds may be set conservatively, so as to minimise use of a reserve compression unit that performs lossy compression (where provided). The data compression unit is also very configurable and the cost/complexity of the bit predictor element **1902** can be traded off against accuracy and performance of the data compression. Additionally it is very scalable and may, for example, be applied only to subsets of the data, e.g. the alpha channel only.

(204) A first example of a lossless compression method which may be implemented in the primary compression unit **202** (shown in FIGS. **2A-C** and described above) and within the data compression unit shown in FIG. **19A** (e.g. within both the bit predictor element **1902** and data compression element **1906** is described in UK patent number 2530312. UK patent number 2530312 describes a data compression system in which 8×8 blocks of image data are compressed by performing colour decorrelation, spatial decorrelation and entropy encoding, and then packing the resulting encoded data values into data packets for storage. In the system described in that patent, a row of data values from a block is processed per iteration. A modified version of Exponential Golomb coding is used for the entropy encoding wherein each coefficient stored a prefix to indicate a range of values that were then identified by a suffix. Having a prefix and suffix for each encoded coefficient makes decoding multiple values simultaneously difficult because it is not simple to quickly determine the positions of the boundaries between encoded bits.

(205) A second example of a lossless compression method which may be implemented in the primary compression unit **202** (shown in FIGS. **2A-C** and described above) can be described with reference to FIGS. **20A-B**, **21A-B** and **22**. As described in UK patent number 2530312 and shown in FIGS. **20A-B**, data compression may comprise the following steps (not necessarily in this order): division into blocks, format conversion (block **2024**), colour channel decorrelation (block **2024**, e.g. as described above), spatial decorrelation (block **2026**), entropy encoding (block **2028**) and then storage.

(206) The entropy encoding (in block **2028**) is performed according to a variable-length coding (VLC) scheme, such that the entropy encoded values will most likely not all have the same number of bits. In general, more probable values are encoded with fewer bits. In this way, it is likely that the total number of bits used to encode all of the data values in a data block will be reduced by the entropy encoding. There are some situations where spatial decorrelation can expand data, and these situations can be treated as special cases (i.e. the data is treated differently to other cases) in order to limit the expansion.

(207) VLC encoding can be reasonably straightforward to perform at a high rate, but VLC decoding at a matching rate can be significantly more difficult because the bit-boundaries between contiguously stored encoded data values are not known until the encoded data is analysed, i.e. the length of the encoded data values is not known until the encoded data is analysed. In particular, when encoding, multiple symbols can be mapped independently to their corresponding codes and code lengths, and then merged into a contiguous output bit stream. However, typically when decoding, each code in the encoded bit stream is examined sequentially in order to determine its length which, in turn, determines the location of the start of the next code. In other words, the bit-boundaries between the different encoded data values need to be found.

(208) In this second example method, entropy encoding is performed on sets of data values (e.g. 2×2 sets of data values). For each set of data values an indication is included in an encoded output (e.g. in a header section thereof) to indicate how many bits are used for each of the encoded data values representing the set of data values. The encoded data values are then included in the encoded output in accordance with the indicated numbers of bits. This system means that the decoding of the encoded data values is simplified (compared to the system of UK patent number 2530312) because a simple read of the indications (e.g. in the header)

allows the bit boundaries between different encoded data values to be determined. In other words the indications allow the sizes of the corresponding encoded data sections to be quickly determined (e.g. in just a single clock cycle). This is described below in more detail with reference to FIGS. 21A-B and 22.

(209) FIG. 21A shows an 8×8 block of data values **2100** which have been spatially decorrelated and remapped by one of the spatial decorrelating modules **2004**. The data values relate to respective pixel values. A method of performing entropy encoding on the block of data values at an entropy encoding module **2006** is described with reference to the flow chart shown in FIG. 22. The entropy encoding module **2006** outputs an encoded data output **2110** (as shown in FIG. 21B) which comprises encoded data values to represent the data values in the 8×8 block **2100**. The encoded data output **2110** has a variable length depending on the compression that can be achieved by the entropy encoding process. In particular, in an example in which each received data value in the block **2100** comprises 8 bits (i.e. 1 byte), the encoded data output **2110** comprises a reference value field **2112** which is 8 bits (1 byte) wide, a size indication field **2114** which is 48 bits (6 bytes) wide, and a variable size field which, in this example has been configured to maintain 8-bit alignment of fields, is anywhere between 8 bits (1 byte) and 448 bits (56 bytes) wide. If the entropy encoding process is not able to losslessly represent the data values in the data block **2100** without exceeding the number of bits available in the encoded data output **2110** (i.e. 504 bits or 63 bytes) then the data values in the data block **2100** are stored in raw form in the encoded data output, with 8 bits per data value. A header section of a data packet in which the encoded data output is placed has an indication to indicate the format of the encoded data within the data packet. For example, the indication may comprise two bits such that up to four different encoding modes can be set, which may include a raw data format, an entropy encoded format and a constant format (in which all of the data values within a block have the same value, such that the value can be stored once (using 8 bits) to represent the value of each of the data values within the block).

(210) The block of data values is received at an entropy encoding module **2006** (block **2202**). The rows and columns of the block are arranged as shown in FIG. 21A. The block of data values **2100** comprises a plurality of sets of data values (most of which in this example are 2×2 sets of data values (relating to 2×2 sets of pixel values), but in other examples, the sets could be different shapes and/or sizes). The block **2100** comprises fifteen 2×2 sets of data values, and one particular set (in the top left of the block) which comprises just three data values. The top left data value (A0) is not entropy encoded, and instead the data value at position A0 is included in the reference value field **2112** of the encoded data output **2110** (block **2204**). Since the top left set represents a similar sized set compared to the other sets in the block, but minus the data value which is to be used as the reference value, the top left set includes fewer (e.g. 3 rather than 4) data values than each of the other sets.

(211) In the example shown in FIG. 21A, the rows in the block have the order A", E", C", G", B", D", F", H", and the columns in the block have the order 0, 4, 2, 6, 1, 3, 5, 7 for the purposes of the entropy encoding. Other examples, may order the rows and columns differently to that shown in FIG. 21A. The example order shown in FIG. 21A does not exactly match the order in which the rows are processed by the spatial decorrelation modules (e.g. row 6 is shown above rows 1 and 3). With the exception of the top left 2×2 set (which only groups three values and is encoded slightly differently), each aligned 2×2 set of values will be encoded together. This example grouping has been chosen such that: (a) the values within a set belong (where possible) to the same spatial frequency, and (b) the values within a set come from pixels that are, where possible, nearby spatially. This is done because these are more likely to be "similar" in magnitude and, hence, the sets can be compressed more effectively.

(212) As described above, the data values are unsigned with a distribution which is biased towards zero (due to the colour correlation and spatial decorrelation processes). Therefore, data values are likely to have leading zeroes. Therefore, the data values can be compressed simply by removing one or more leading zeroes from the data values (where possible). An indication is used to indicate how many leading zeroes have been removed.

(213) The top left set of four data values are treated differently in that: (a) the top left pixel is used as the reference and stored separately, and (b) the remaining three values have a different encoding scheme that has been "trained" on a large set of image data so that: (i) for alignment purposes, the total number of bits used to represent the three values is a multiple of 4 (this matches the 2×2 blocks), and (ii) the total storage cost (for the training image set) has been minimised. The particular set of data values in the top left of the block is processed (block **2206**). Specifically, an indication for the particular set is included in the size indication field **2114** (block **2208**). Each indication in the size indication field has three bits, for a respective set of data values in the block **2100**. This 3-bit indication can be used as an index into a Look Up Table to retrieve a number of bits for each data value in the particular set, such that the number of bits for each value is sufficient to store that value. There are sixteen sets of data values in the block **2100**, such that the size indication field comprises 48 bits (i.e. 3×16 bits), or 6 bytes. The fifteen sets of four data values in a block will be encoded with numbers of bits which are multiples of four (because, as described below, each encoded data value in a set of four data values has the same number of bits), and on any given row each of these fifteen sets contribute a multiple of two bits. To maintain a convenient alignment of compressed data (as restricting alignment of data can reduce hardware complexity) restrictions are applied to the available choices for the top-left set which only has 3 elements. The top left set has three data values, one of which (E0) is in row 0. For hardware efficiency, it is desirable that the total data for a row is an even number of bits, so the possible lengths of the encoded data values which can be used to represent this data value are restricted to even numbers. Similarly, the combined lengths of the two encoded data values which are used to represent the other two data values of this set (A4 and E4) sum to an even value. Furthermore, in this example, the combined length of all the encoded data values in the encoded data output will be rounded up to the nearest byte (so that the encoded data outputs are aligned on byte boundaries), and all of the other sets of data values in the block. Therefore, the total number of bits used to encode the top left set of data values is also a multiple of four.

(214) With a size indicator including three bits, eight sets of lengths for the encoded data values can be set. For example, the table below shows some possible lengths of the encoded data values which may be represented by the eight different size indications for the top left block. The lengths shown in the table below were made by evaluating a very large set of images and choosing the combinations (from a very large set of possible combinations) that resulted in the lowest overall storage costs.

(215) TABLE-US-00023 Size indication E0 A4 E4 Total Length 000 0 0 0 0 001 2 1 1 4 010 2 3 3 8 011 4 3 5 12 100 4 5 3 12 101 6 5 5 16 110 6 7 7 20 111 8 8 8 24

(216) The entropy encoding module **2006** chooses the coding from the possible options (e.g. as shown in the table above) with the least total length that is able to (losslessly) represent the {E0,A4,E4} triple by removing leading zeroes from the data values. In the event of a tie (e.g. if deciding between size indications 011 and 100 in the example shown in the table above) either tied code could be chosen, but as an example, the code with the least numerical encoding may be chosen (e.g. 011).

(217) The encoded data values for the top left set are included in the variable size field **2116** (e.g. at the start of the variable size field **2116**) of the encoded data output **2110** (block **2210**).

(218) The remaining (e.g. fifteen) sets of data values in the block **2100** are then processed (block **2212**). Specifically, a number of bits (e.g. a minimum number of bits) for representing the maximum value of the data values of a set is determined (block **2214**), and an indication of the determined number of bits is included in the size indication field **2114** of the encoded data output **2110** for the set (block **2216**). Block **2214** may be implemented by finding which of the data values in the set has the smallest number of leading zeros, and then by identifying the smallest available encoding that will suffice. It is noted that not all possible data lengths may be available to be encoded. For example, as described in the example below, a data length of seven bits is not available to be encoded. Therefore, the determined number of bits may be a minimum “valid” number of bits which can be used to represent the maximum value of the received data values in the set, wherein a number of bits is valid if it can be indicated by an indication to be included in the size indication field **2114**. For example, there may be a predetermined set of valid numbers of bits which can be indicated by said indication (e.g. as listed in the table above), and the minimum valid number of bits may be the minimum of the valid numbers of bits in the predetermined set which is sufficient to represent the maximum value of the received data values in the set.

(219) When the number of bits which is to be used to represent each of the encoded data values in a set has been determined, some leading zeroes can be removed from each of the data values in the set (if appropriate) to thereby determine the encoded data values, such that each of the encoded data values in a set has the determined number of bits for that set.

(220) The encoded data values representing the data values in the set are included in the variable size field **2116** of the encoded data output **2110** (block **2218**). The order in which the encoded data values are included in the variable size field **2116** is predetermined and corresponds with the order in which the indications are included in the size indication field **2114**, such that when the encoded data values are decoded, the positions of the decoded data values within the block are known.

(221) For example, if the four data values in a set are 00000011, 00000100, 00000001 and 00000000, then 00000100 is the maximum value in the set, and three is the minimum number of bits that can be used to represent the maximum value. Three is a valid number of bits in the example shown in the table above in the sense that an indication can be included to indicate that each of the encoded data values of the set have three bits. Therefore, five leading zeroes can be removed from each of the data values in the set. Therefore, the encoded data values for the set are 011, 100, 001 and 000, and these encoded data values are included in the variable size field **2116**. An indication (e.g. 011) to indicate that three bits are used to represent each encoded data value in the set is included in the size indication field **2114**. Different sets (e.g. 2×2 sets) can use different numbers of bits for their encoded data values, and each set has its own indication in the size indication field **1304** to indicate how many bits are used in the variable size field **2116** for the encoded data values of that set.

(222) As an example, the table below shows how the indications may correspond to numbers of bits for the 2×2 sets.

(223) TABLE-US-00024 Size indication Number of bits per encoded data value 000 0 001 1 010 2 011 3 100 4 101 5 110 6 111 8

(224) It is noted that there is no indication to indicate that seven bits are used for each encoded data value. Therefore, in this example, if the maximum 8-bit value within a set has just one leading zero then the minimum valid number of bits which can be used to represent each of the encoded data values in that set is eight (not seven). This omission of seven from the predetermined set of valid lengths for encoded data values was chosen by examining all the possible choices of number of 3 bits to a set of 8 choices. Since there are 9 options (bit lengths 0 to 8), one option is left out. Having evaluated all the possibilities (e.g. trialing leaving out 0, or 1, or 2 . . .) against a large test suite of images, it was confirmed that leaving out “7” gives the best level of compression. It is noted that the likelihood of data values having only one leading zero is small because the colour decorrelation and the spatial decorrelation processes cause the data values to be biased towards zero.

(225) It should be apparent that the examples described herein relate to 8-bit data values, but in other examples, the data values may include different numbers of bits (e.g. 6-bit data values), and a person skilled in the art would be able to modify the disclosure provided herein to operate with these different numbers of bits.

(226) For example, if each of the data values received from a spatial decorrelation module has n bits, and the determined minimum valid number of bits for a particular set of received data values is m, where m≤n, then each of the received data values of the particular set has at least (n−m) leading zeroes. In this case, each of the encoded data values for the particular set consists of the m least significant bits of a corresponding n-bit received data value of the particular set. Referring to the same example as above, if the four n-bit data values (where n=8) in a set are 00000011, 00000100, 00000001 and 00000000, then the determined minimum valid number of bits for representing the maximum value of the set, m=3. A data length of 3-bits per encoded data value is a valid encoding from the table above. Each of the data values has at least five leading zeroes. Therefore, the m-bit encoded data values for the set are 011, 100, 001 and 000.

(227) Different sets within the block **2100** can be processed in parallel by an entropy encoding module **2006**. For example, the indications for the different sets within the block **2100** can be determined and included in the size indication field **2114** in parallel operations within the entropy encoding module **2006**. However, since the encoded data values can have variable lengths, they are included in the variable size field **2116** in a predetermined sequence.

(228) When all of the encoded data values of the block have been included in the encoded data output **2110** then, the encoded data output is outputted from the entropy encoding module **2006** (block **2220**). The encoded data output representing a block of pixel values is provided to the packing module **2008** for each of the colour channels, where they are packed together. For example, the packing module **2008** places the encoded data outputs for the block of data values from the different colour channels into a data packet. In other words, a plurality of encoded data outputs are formed for a respective plurality of colour channels relating to the same pixels, and the plurality of encoded data outputs for a group of pixels (e.g. an 8×8 block) are packed together into a data

packet for storage. The encoded data block can then be sent for storage in the memory **106**, e.g. via the memory interface **114** (as shown in FIG. **1**).

(229) A header is stored with a group of encoded data blocks, e.g. the data blocks representing an image or a frame. For each encoded data block (e.g. each 8×8 block of encoded data values), a header indication is included in the header to indicate the size of the compressed data for the respective encoded data block. The header is stored separately from the encoded data values (e.g. in a dedicated portion of the encoded data output), and due to its small size, a significant portion of the header data may remain resident in a cache within the GPU. The header allows the size of the encoded data block to be known before the encoded data block is retrieved from memory, so an encoded data block can be retrieved without retrieving extraneous data. This is very useful for allowing random access to different parts of encoded data. In general, any number of bits may be used in the header as a header indication for a respective encoded data block, but in a specific example described below, four bits are used for each of the header indications.

(230) As noted above, whilst the method described above with reference to FIGS. **20A-B**, **21A-B** and **22** and the method described in UK patent number 2530312 are both inherently lossless, if either is used in combination with a lossy pre-processing step (e.g. as described above with reference to FIG. **14** or with reference to FIGS. **16** and **18A** when one or more bits from the fourth pseudo-channel are discarded), the resultant compressed data is no longer lossless.

(231) Any reference to particular logic gates above (e.g. OR, AND gates, etc.) is by way of example only and it will be appreciated that they may be replaced by logic elements that provide the same functionality and may be more broadly referred to as logic blocks.

(232) The description above describes many different methods, including methods of data compression (e.g. FIGS. **2A-C**, **3A**, **4-11** and **19A**), methods of bit mapping (e.g. FIGS. **12A-12B**), including conversion from 10-bit to 8-bit format (e.g. FIGS. **14** and **16**). These methods may be used individually or two or more of these methods may be used together in various combinations and several examples are described above. Also described above are the reverse operations (e.g. data compression, inverse mapping and conversion from 8-bit to 10-bit format) and the hardware for implementing any of these methods or any combination of these methods.

(233) FIG. **23** shows a computer system in which the data compression and decompression methods and apparatus described herein may be implemented. The computer system comprises a CPU **2302**, a GPU **2304**, a memory **2306** and other devices **2314**, such as a display **2316**, speakers **2318** and a camera **2320**. A data compression and/or decompression block **2321** (which may implement any of the methods described herein) is implemented on the GPU **2304**. In other examples, the data compression and/or decompression block **2321** may be implemented on the CPU **2302**. The components of the computer system can communicate with each other via a communications bus **2322**.

(234) The data compression hardware of FIGS. **2A-C**, **12B** and **13B** are shown as comprising a number of functional blocks. This is schematic only and is not intended to define a strict division between different logic elements of such entities. Each functional block may be provided in any suitable manner. It is to be understood that intermediate values described herein as being formed by data compression hardware need not be physically generated by the data compression hardware at any point and may merely represent logical values which conveniently describe the processing performed by the data compression hardware between its input and output.

(235) The data compression and decompression hardware described herein (including any hardware that is arranged to implement any of the methods described above) may be embodied in hardware on an integrated circuit. The data compression and decompression hardware described herein may be configured to perform any of the methods described herein. Generally, any of the functions, methods, techniques or components described above can be implemented in software, firmware, hardware (e.g., fixed logic circuitry), or any combination thereof. The terms “module,” “functionality,” “component,” “element,” “unit,” “block” and “logic” may be used herein to generally represent software, firmware, hardware, or any combination thereof. In the case of a software implementation, the module, functionality, component, element, unit, block or logic represents program code that performs the specified tasks when executed on a processor. The algorithms and methods described herein could be performed by one or more processors executing code that causes the processor(s) to perform the algorithms/methods. Examples of a computer-readable storage medium include a random-access memory (RAM), read-only memory (ROM), an optical disc, flash memory, hard disk memory, and other memory devices that may use magnetic, optical, and other techniques to store instructions or other data and that can be accessed by a machine.

(236) The terms computer program code and computer readable instructions as used herein refer to any kind of executable code for processors, including code expressed in a machine language, an interpreted language or a scripting language. Executable code includes binary code, machine code, bytecode, code defining an integrated circuit (such as a hardware description language or netlist), and code expressed in a programming language code such as C, Java or OpenCL. Executable code may be, for example, any kind of software, firmware, script, module or library which, when suitably executed, processed, interpreted, compiled, executed at a virtual machine or other software environment, cause a processor of the computer system at which the executable code is supported to perform the tasks specified by the code.

(237) A processor, computer, or computer system may be any kind of device, machine or dedicated circuit, or collection or portion thereof, with processing capability such that it can execute instructions. A processor may be any kind of general purpose or dedicated processor, such as a CPU, GPU, System-on-chip, state machine, media processor, an application-specific integrated circuit (ASIC), a programmable logic array, a field-programmable gate array (FPGA), physics processing units (PPUs), radio processing units (RPU), digital signal processors (DSPs), general purpose processors (e.g. a general purpose GPU), microprocessors, any processing unit which is designed to accelerate tasks outside of a CPU, etc. A computer or computer system may comprise one or more processors. Those skilled in the art will realize that such processing capabilities are incorporated into many different devices and therefore the term ‘computer’ includes set top boxes, media players, digital radios, PCs, servers, mobile telephones, personal digital assistants and many other devices.

(238) It is also intended to encompass software which defines a configuration of hardware as described herein, such as HDL

(hardware description language) software, as is used for designing integrated circuits, or for configuring programmable chips, to carry out desired functions. That is, there may be provided a computer readable storage medium having encoded thereon computer readable program code in the form of an integrated circuit definition dataset that when processed (i.e. run) in an integrated circuit manufacturing system configures the system to manufacture data compression and/or decompression hardware configured to perform any of the methods described herein, or to manufacture data compression and/or decompression hardware comprising any apparatus described herein. An integrated circuit definition dataset may be, for example, an integrated circuit description.

(239) Therefore, there may be provided a method of manufacturing, at an integrated circuit manufacturing system, data compression and/or decompression hardware as described herein. Furthermore, there may be provided an integrated circuit definition dataset that, when processed in an integrated circuit manufacturing system, causes the method of manufacturing data compression and/or decompression hardware to be performed.

(240) An integrated circuit definition dataset may be in the form of computer code, for example as a netlist, code for configuring a programmable chip, as a hardware description language defining an integrated circuit at any level, including as register transfer level (RTL) code, as high-level circuit representations such as Verilog or VHDL, and as low-level circuit representations such as OASIS (RTM) and GDSII. Higher level representations which logically define an integrated circuit (such as RTL) may be processed at a computer system configured for generating a manufacturing definition of an integrated circuit in the context of a software environment comprising definitions of circuit elements and rules for combining those elements in order to generate the manufacturing definition of an integrated circuit so defined by the representation. As is typically the case with software executing at a computer system so as to define a machine, one or more intermediate user steps (e.g. providing commands, variables etc.) may be required in order for a computer system configured for generating a manufacturing definition of an integrated circuit to execute code defining an integrated circuit so as to generate the manufacturing definition of that integrated circuit.

(241) An example of processing an integrated circuit definition dataset at an integrated circuit manufacturing system so as to configure the system to manufacture data compression and/or decompression hardware will now be described with respect to FIG. 24.

(242) FIG. 24 shows an example of an integrated circuit (IC) manufacturing system **2402** which is configured to manufacture data compression and/or decompression hardware as described in any of the examples herein. In particular, the IC manufacturing system **2402** comprises a layout processing system **2404** and an integrated circuit generation system **2406**. The IC manufacturing system **2402** is configured to receive an IC definition dataset (e.g. defining data compression and/or decompression hardware as described in any of the examples herein), process the IC definition dataset, and generate an IC according to the IC definition dataset (e.g. which embodies data compression and/or decompression hardware as described in any of the examples herein). The processing of the IC definition dataset configures the IC manufacturing system **2402** to manufacture an integrated circuit embodying data compression and/or decompression hardware as described in any of the examples herein.

(243) The layout processing system **2404** is configured to receive and process the IC definition dataset to determine a circuit layout. Methods of determining a circuit layout from an IC definition dataset are known in the art, and for example may involve synthesising RTL code to determine a gate level representation of a circuit to be generated, e.g. in terms of logical components (e.g. NAND, NOR, AND, OR, MUX and FLIP-FLOP components). A circuit layout can be determined from the gate level representation of the circuit by determining positional information for the logical components. This may be done automatically or with user involvement in order to optimise the circuit layout. When the layout processing system **2404** has determined the circuit layout it may output a circuit layout definition to the IC generation system **2406**. A circuit layout definition may be, for example, a circuit layout description.

(244) The IC generation system **2406** generates an IC according to the circuit layout definition, as is known in the art. For example, the IC generation system **2406** may implement a semiconductor device fabrication process to generate the IC, which may involve a multiple-step sequence of photo lithographic and chemical processing steps during which electronic circuits are gradually created on a wafer made of semiconducting material. The circuit layout definition may be in the form of a mask which can be used in a lithographic process for generating an IC according to the circuit definition. Alternatively, the circuit layout definition provided to the IC generation system **2406** may be in the form of computer-readable code which the IC generation system **2406** can use to form a suitable mask for use in generating an IC.

(245) The different processes performed by the IC manufacturing system **2402** may be implemented all in one location, e.g. by one party. Alternatively, the IC manufacturing system **2402** may be a distributed system such that some of the processes may be performed at different locations, and may be performed by different parties. For example, some of the stages of: (i) synthesising RTL code representing the IC definition dataset to form a gate level representation of a circuit to be generated, (ii) generating a circuit layout based on the gate level representation, (iii) forming a mask in accordance with the circuit layout, and (iv) fabricating an integrated circuit using the mask, may be performed in different locations and/or by different parties.

(246) In other examples, processing of the integrated circuit definition dataset at an integrated circuit manufacturing system may configure the system to manufacture data compression and/or decompression hardware without the IC definition dataset being processed so as to determine a circuit layout. For instance, an integrated circuit definition dataset may define the configuration of a reconfigurable processor, such as an FPGA, and the processing of that dataset may configure an IC manufacturing system to generate a reconfigurable processor having that defined configuration (e.g. by loading configuration data to the FPGA).

(247) In some embodiments, an integrated circuit manufacturing definition dataset, when processed in an integrated circuit manufacturing system, may cause an integrated circuit manufacturing system to generate a device as described herein. For example, the configuration of an integrated circuit manufacturing system in the manner described above with respect to FIG. 24 by an integrated circuit manufacturing definition dataset may cause a device as described herein to be manufactured.

(248) In some examples, an integrated circuit definition dataset could include software which runs on hardware defined at the dataset or in combination with hardware defined at the dataset. In the example shown in FIG. 24, the IC generation system may further be configured by an integrated circuit definition dataset to, on manufacturing an integrated circuit, load firmware onto that

integrated circuit in accordance with program code defined at the integrated circuit definition dataset or otherwise provide program code with the integrated circuit for use with the integrated circuit.

(249) Those skilled in the art will realize that storage devices utilized to store program instructions can be distributed across a network. For example, a remote computer may store an example of the process described as software. A local or terminal computer may access the remote computer and download a part or all of the software to run the program. Alternatively, the local computer may download pieces of the software as needed, or execute some software instructions at the local terminal and some at the remote computer (or computer network). Those skilled in the art will also realize that by utilizing conventional techniques known to those skilled in the art that all, or a portion of the software instructions may be carried out by a dedicated circuit, such as a DSP, programmable logic array, or the like.

(250) The applicant hereby discloses in isolation each individual feature described herein and any combination of two or more such features, to the extent that such features or combinations are capable of being carried out based on the present specification as a whole in the light of the common general knowledge of a person skilled in the art, irrespective of whether such features or combinations of features solve any problems disclosed herein. In view of the foregoing description it will be evident to a person skilled in the art that various modifications may be made within the scope of the invention.

Claims

1. A method of converting 10-bit format pixel data into 8-bit format data, the method comprising, for a pixel: generating an error value for a 10-bit input data channel by comparing a first number of least significant bits (LSBs) of the input data channel and a same first number of bits derived from a truncated version of the channel; and outputting an 8-bit channel comprising the error value.
2. The method according to claim 1, wherein the truncated version of the channel comprises 7 bits.
3. The method according to claim 1, wherein generating an error value for the 10-bit input data channel by comparing a first number of LSBs of the input data channel and a same first number of bits derived from the truncated version of the channel comprises: generating an error value for the 10-bit input data channel by combining using an XOR logic block three LSBs of the input data channel and three bits derived from the truncated version of the channel.
4. The method according to claim 1, wherein the first number of bits derived from the truncated version of the channel comprises the three most significant bits (MSBs) of the truncated version of the channel.
5. The method according to claim 1, further comprising: prior to generating the error values, expanding, using bit replication, the truncated version of the channels to 10 bits; and wherein the first number of bits derived from the truncated version of the channel comprises the three LSBs of the expanded truncated version of the channel.
6. The method according to claim 1, wherein the truncated version of the channel is encoded using Gray coding.
7. The method according to claim 1, further comprising outputting an 8-bit channel comprising the truncated version of the channel with a different appended bit from a further input data channel appended as an MSB.
8. The method according to claim 1, wherein the 10-bit pixel data comprises RGBA pixel data.
9. A method of compressing 10-bit format pixel data comprising: converting 10-bit format pixel data into 8-bit format data using the method as set forth in claim 1; and compressing a block of 8-bit format data using a lossless compression method to generate a first compressed data block.
10. The method according to claim 9, further comprising: assessing the first compressed data block against a target compression threshold; in response to determining that the first compressed data block satisfies the target compression threshold, outputting the first compressed data block; and in response to determining that the first compressed data block does not satisfy the target compression threshold, discarding one or more bits from the 8-bit channel comprising the error values in the block of 8-bit format data to generate a reduced size block of data and compressing the reduced size block of data using the lossless compression method to generate a second compressed data block, wherein the second compressed data block is smaller than the first compressed data block.
11. The method according to claim 10, further comprising: outputting the second compressed data block.
12. The method according to claim 10, further comprising: assessing the second compressed data block against the target compression threshold; in response to determining that the second compressed data block satisfies the target compression threshold, outputting the second compressed data block; and in response to determining that the second compressed data block does not satisfy the target compression threshold, outputting, instead of the second compressed data block, a third compressed data block generated from the block of 8-bit format data, wherein the third compressed data block is generated using a lossy compression technique, and wherein the lossy compression technique is arranged to ensure that the third compressed data block satisfies the target compression threshold.
13. A method of converting 8-bit format pixel data into 10-bit format data, the method comprising: using bit replication to increase an input data channel to a 10-bit channel; and for the channel, modifying a first number of least significant bits (LSBs) by combining them with corresponding bits from an 8-bit channel to generate a corresponding 10-bit output channel.
14. The method according to claim 13, further comprising: Gray decoding the 2-bit last output channel.
15. A data decompression unit arranged to convert 8-bit format pixel data into 10-bit format data, the hardware comprising: bit replication logic arranged, for a pixel, to use bit replication to increase an input data channel to a 10-bit channel; and hardware logic arranged, for the channel, to modify a first number of least significant bits (LSBs) by combining them with corresponding bits from an 8-bit channel to generate a corresponding 10-bit output channel.
16. The data decompression unit of claim 15, wherein the data compression unit is embodied in hardware on an integrated circuit.
17. A non-transitory computer readable storage medium having stored thereon computer executable code that when executed causes at least one processor to perform the method as set forth in claim 1.
18. A non-transitory computer readable storage medium having stored thereon a computer readable dataset description of an

integrated circuit that, when processed in an integrated circuit manufacturing system, causes the integrated circuit manufacturing system to manufacture a data decompression unit as set forth in claim 15.

19. An integrated circuit manufacturing system comprising: a computer readable storage medium having stored thereon a non-transitory computer readable dataset description of an integrated circuit that describes a data decompression unit as set forth in claim 15; a layout processing system configured to process the integrated circuit description so as to generate a circuit layout description of an integrated circuit embodying the data compression unit; and an integrated circuit generation system configured to manufacture the data compression unit according to the circuit layout description.
