

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12418687
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Katsavounidis; Ioannis

Optimizing encoding operations when generating a buffer-constrained version of a media title

Abstract

In various embodiments, a buffer-based encoding application generates a first convex hull of subsequence encode points based on multiple encoding points and a first subsequence associated with a media title. The buffer-based encoding application then generates a first global convex hull of media encode points based on a transmission buffer constraint, the first convex hull, and a second global convex hull of media encode points. Notably, the second global convex hull is associated with a portion of the media title that occurs before the first subsequence in a playback order for the media title. Subsequently, the subsequence-based encoding application selects a first media encode point included in the first global convex hull based on a media metric and determines a first encoded media sequence based on the selected media encode point. The first encoded media sequence is subsequently streamed to an endpoint device during playback of the media title.

Inventors:	Katsavounidis; Ioannis (San Jose, CA)
Applicant:	NETFLIX, INC. (Los Gatos, CA)
Family ID:	1000008820651
Assignee:	NETFLIX, INC. (Los Gatos, CA)
Appl. No.:	17/164548
Filed:	February 01, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20210160550 A1	May. 27, 2021

Related U.S. Application Data

continuation parent-doc US 16244033 20190109 US 10911791 child-doc US 17164548

Publication Classification

Int. Cl.: H04N21/231 (20110101); H04N21/235 (20110101); H04N21/236 (20110101);
H04N21/2662 (20110101)

U.S. Cl.:

CPC H04N21/231 (20130101); H04N21/235 (20130101); H04N21/236 (20130101);
H04N21/2662 (20130101);

Field of Classification Search

CPC: H04N (21/231); H04N (21/235); H04N (21/236); H04N (21/2662)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8837601	12/2013	Ronca et al.	N/A	N/A
2004/0165783	12/2003	Reynolds et al.	N/A	N/A
2007/0002946	12/2006	Bouton et al.	N/A	N/A
2011/0019974	12/2010	Konoshima	386/E5.028	G11B 27/036
2011/0052087	12/2010	Mukherjee	N/A	N/A
2011/0090949	12/2010	Gu	N/A	N/A
2013/0223510	12/2012	Coudurier et al.	N/A	N/A
2013/0329781	12/2012	Su et al.	N/A	N/A
2016/0254976	12/2015	Hayashi	370/252	H04L 43/0894
2016/0295254	12/2015	Chen et al.	N/A	N/A
2017/0064342	12/2016	Botsford	N/A	H04N 21/2381
2017/0264820	12/2016	Shibayama	N/A	H04N 19/15
2018/0063549	12/2017	Amer et al.	N/A	N/A
2018/0302456	12/2017	Katsavounidis et al.	N/A	N/A
2018/0343458	12/2017	Katsavounidis et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
103621085	12/2013	CN	N/A
107145888	12/2016	CN	N/A
108965930	12/2017	CN	N/A
WO-2005029868	12/2004	WO	H04N 19/102
2018/156997	12/2017	WO	N/A

OTHER PUBLICATIONS

Jan De Cock et al., Complexity-based consistent-quality encoding in the cloud, Sep. 1, 2016, 2016 IEEE International Conference on Image Processing (ICIP) (pp. 1484-1488) (Year: 2016). cited by examiner

Complexity-Based Consistent-Quality Encoding in the Cloud (Year: 2016). cited by examiner

Zhang et al., “Rate control optimization for live streaming using fixed-lime playback buffer”, Soft Computing, DOI:10.1007/S00500-015-1854-7, vol. 21, No. 5, Sep. 23, 2015, pp. 1219-1228. cited by applicant

Wiegand et al., “Overview of the H.264/AVC Video Coding Standard”, DOI:10.1109/TCSVT.2003.815165, IEEE Transactions on Circuits and Systems for Video Technology, vol. 13, No. 7, Jul. 2003, pp. 560-576. cited by applicant

“The Netflix Tech Blog: High Quality Video Encoding at Scale”, Retrieved on May 19, 2020, <https://netflixtechblog.com/high-quality-video-encoding-at-scale-d159db052746>, Dec. 9, 2015, 6 pages. cited by applicant

Ortega et al., “Rate-Distortion Methods for Image and Video Compression”, IEEE Signal Processing Magazine, vol. 15, No. 6, Nov. 1998, pp. 23-50. cited by applicant

Richardson, Iain E.G., “H. 264 and MPEG-4 Video Compression Video Coding for Next-generation Multimedia”, John Wiley & Sons, 2003, 307 pages. cited by applicant

Keys, Robert G., “Cubic Convolution Interpolation for Digital Image Processing”, IEEE Transactions on Acoustics, Speech, and Signal Processing, vol. ASSP-29, No. 6, Dec. 1981, pp. 1153-1160. cited by applicant

“Lanczos resampling—Wikipedia”, Retrieved on May 19, 2020, http://en.wikipedia.org/wiki/Lanczos_resampling, 5 pages. cited by applicant

“The Netflix Tech Blog: Per-Title Encode Optimization”, Retrieved on May 19, 2020, <https://netflixtechblog.com/per-title-encode-optimization-7e99442b62a2>, Dec. 14, 2015, 13 pages. cited by applicant

“The Netflix Tech Blog: Toward A Practical Perceptual Video Quality Metric”, Retrieved on May 19, 2020, <https://netflixtechblog.com/toward-a-practical-perceptual-video-quality-metric-653f208b9652>, Jun. 6, 2016, 21 pages. cited by applicant

“X264—Wikipedia”, Retrieved on May 19, 2020, <https://en.wikipedia.org/wiki/X264>, 5 pages. cited by applicant

Tavakoli et al., “Quality of experience of adaptive video streaming: Investigation in service parameters and subjective quality assessment methodology”, doi.org/10.1016/j.image.2015.05.001, Signal Processing: Image Communication, vol. 39, 2015, pp. 1-12. cited by applicant

Katsavounidis, Ioannis, “Netflix-“El Fuente” video sequence details and scenes”, http://www.cdv1.org/documents/ElFuente_summary.pdf, Jul. 28, 2015, 64 pages. cited by applicant

Sullivan et al., “Rate-Distortion Optimization for Video Compression”, IEEE Signal Processing Magazine, vol. 15, No. 6, Nov. 1998, pp. 74-90. cited by applicant

Thiede et al., “PEAQ—The ITU Standard for Objective Measurement of Perceived Audio Quality”, Journal of the Audio Engineering Society, vol. 48, No. 1/2, Jan./Feb. 2000, pp. 3-29. cited by applicant

The Consumer Digital Video Library, Retrieved on May 19, 2020, <http://www.cdv1.org>, 1 page. cited by applicant

Katsavounidis et al., “Native resolution detection of video sequences,” in SMPTE 2015 Annual Technical Conference and Exhibition, SMPTE, 2015, <https://ieeexplore.ieee.org/abstract/document/7399661>. Manuscript, 26 pages. cited by applicant

Oogami et al., “A study on a compliant VBV Buffer Model for Distributed Video Transcoding”, Information Processing Society of Japan Report-of-Research IPSJ SIG Technical Reports, vol. 2004, Mar. 4, 2004, pp. 31-36. cited by applicant

Mehrotra et al., “Rate-Distortion Optimized Client Side Rate Control for Adaptive Media

Primary Examiner: Abouzahra; Hesham K

Attorney, Agent or Firm: Artega Law Group, LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of the co-pending U.S. patent application titled, “OPTIMIZING ENCODING OPERATIONS WHEN GENERATING A BUFFER-CONSTRAINED VERSION OF A MEDIA TITLE,” filed on Jan. 9, 2019 and having U.S. Ser. No. 16/244,033. The subject matter of this related application is hereby incorporated herein by reference.

BACKGROUND

Field of the Various Embodiments

(1) Embodiments of the present invention relate generally to encoding technology and, more specifically, to optimizing encoding operations when generating a buffer-constrained version of a media title.

Description of the Related Art

(2) A typical media streaming service provides access to a library of media titles that can be viewed on a range of different endpoint devices. Each endpoint device may connect to the media streaming service under different connection conditions that affect bandwidth and latency. In addition, each different endpoint device may include different hardware for outputting the media title to end user(s). For example, a given endpoint device could include a display screen having a particular screen size and a particular screen resolution.

(3) In many implementations, an endpoint device that connects to a media streaming service executes an endpoint application that determines, for a given media title, an appropriate encoded version of the media title to stream to the endpoint device based on the connection conditions and the properties of the endpoint device. More specifically, the endpoint application attempts to select a particular encoded version of the media title that provides the best possible visual quality during playback of the media title on the endpoint device while avoiding playback interruptions due to buffering or re-buffering.

(4) In some implementations, the endpoint application selects the particular encoded version of the media title based on a bitrate ladder. The bitrate ladder is designed to achieve a target visual quality during playback of a media title based on an available bandwidth. Each rung in the bitrate ladder specifies a different bitrate-resolution pair corresponding to a different pre-generated encoded version of the media title. To generate the encoded version of a media title corresponding to a given bitrate-resolution pair, the media content associated with the media title is sampled at the resolution to generate sampled media content. A rate control value is selected based on the bitrate and then an encoder is used to encode the sampled media content.

(5) One drawback of the above “monolithic” encoding technique is that the complexity of the media content associated with a given media title oftentimes varies across the media title; whereas, the resolution and the rate control value used to encode the media content do not vary. As a result, encoding relatively simple portions of the media title may consume more computational and storage resources than what is necessary to meet the target visual quality. For example, a relatively simple portion of a media title could have the same visual quality regardless of whether that portion of media content is encoded using a bitrate of 560 kilobits per second (kbps) or using a bitrate of

3000 kbps. Among other things, such encoding inefficiencies waste computational and storage resources and increase the bandwidth required to stream encoded versions of media titles to endpoint devices.

(6) In other implementations, to reduce these types of encoding inefficiencies, a media streaming service provider varies the resolution and the rate control value across the media title. Typically, a subsequence-based encoding application partitions the media title into different subsequences or sets of frames which are characterized by similar properties. The subsequence-based encoding application then encodes each subsequence numerous times at a variety of different resolutions and rate control values to generate encoded subsequences. Subsequently, the subsequence-based encoding application performs optimization operations to generate different optimized encoded versions of the media title. Each optimized encoded version of the media title includes a different combination of the encoded subsequences that span the length of the media title, and the resolution and/or rate control value may vary between the constituent encoded subsequences.

(7) One drawback of the above “subsequence-based” encoding technique is that the optimized encoded versions of a media title do not necessarily comply with the transmission buffer constraints that are oftentimes associated with transmitting to and buffering at the endpoint devices. In general, if an encoded version of a media title complies with the relevant transmission buffer constraints, then the endpoint device is going to be more able to properly buffer and play back the media title, in particular on a constant bitrate transmission channel. If, however, the encoded version of a media title does not comply with the relevant transmission buffer constraints, then the endpoint device is going to be less able to properly buffer and play back the media title. In the latter case, playback interruptions that reduce the overall quality of the viewing experience can occur when streaming the media title to the endpoint device.

(8) As an example, an encoded version of a media title that is transmitted through a constant bandwidth channel, such as those used in broadcast television, is typically required to comply with a constant bitrate (CBR) Video Buffer Verifier (VBV) constraint. In a constant bandwidth channel, encoded bits arrive at a decoder associated with an endpoint device at a fixed rate and are stored in a transmission buffer having a fixed size. As part of displaying a particular frame of the media title, the decoder removes the associated encoded bits from the transmission buffer. Because the number of bits used to encode each frame can vary across the media title, the number of encoded bits that the transmission buffer stores can also vary during playback of the media title. If an encoded version of a media title is compliant with the CBR VBV constraint, then the number of bits used to encode each frame is limited in an effort to ensure that the transmission buffer does not overflow or underflow during playback of the media title. However, if the transmission buffer overflows, then the transmission buffer may not have space to store the incoming encoded bits. Consequently, frame drop can occur. Further, if the transmission buffer underflows, then no encoded bits can be retrieved from the transmission buffer, and frame freeze can occur.

(9) As the foregoing illustrates, what is needed in the art are more effective techniques for encoding media titles.

SUMMARY

(10) One embodiment of the present invention sets forth a computer-implemented method for encoding a media title. The method includes generating a first convex hull of subsequence encode points based on a first set of encoding points and a first subsequence associated with a media title; generating a first global convex hull of media encode points based on a transmission buffer constraint, the first convex hull, and a second global convex hull of media encode points associated with a portion of the media title that occurs before the first subsequence in a playback order for the media title; selecting a first media encode point included in the first global convex hull based on a media metric; and determining a first encoded media sequence based on the first media encode point, wherein the first encoded media sequence is subsequently streamed to an endpoint device during playback of the media title.

(11) At least one technical advantage of the disclosed techniques relative to the prior art is that encoded media sequences are generated based on sets of individually encoded subsequences and transmission buffer constraints. Notably, the individually encoded subsequences included in each encoded media sequence can be associated with different encoding points. Consequently, encoding inefficiencies typically associated with prior art monolithic encoding techniques are reduced. Further, by ensuring that the optimized encoded media sequences are compliant with the transmission buffer constraints, degraded viewing experiences attributable to transmission and buffering limitations when streaming the media title to endpoint devices are reduced. These technical advantages provide one or more technological advancements over prior art approaches.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) So that the manner in which the above recited features of the various embodiments can be understood in detail, a more particular description of the inventive concepts, briefly summarized above, may be had by reference to various embodiments, some of which are illustrated in the appended drawings. It is to be noted, however, that the appended drawings illustrate only typical embodiments of the inventive concepts and are therefore not to be considered limiting of scope in any way, and that there are other equally effective embodiments.
- (2) FIG. 1 is a conceptual illustration of a system configured to implement one or more aspects of the present invention;
- (3) FIG. 2 is a more detailed illustration of the buffer-based encoding application of FIG. 1, according to various embodiments of the present invention;
- (4) FIG. 3 is a more detailed illustration of the encoding point list of FIG. 2, according to various embodiments of the present invention;
- (5) FIG. 4 is a more detailed illustration of one of the subsequence encode sets of FIG. 2, according to various embodiments of the present invention;
- (6) FIG. 5 illustrates an exemplary convex hull that is generated by the convex hull generator of FIG. 2, according to various embodiments of the present invention;
- (7) FIG. 6 is a more detailed illustration of one of the media encode points of FIG. 2, according to various embodiments of the present invention;
- (8) FIG. 7 is a more detailed illustration of the subsequence iterator of FIG. 2, according to various embodiments of the present invention;
- (9) FIG. 8 illustrates an exemplary global convex hull and an exemplary optimized encoded media sequence that are generated by the buffer-based encoding application of FIG. 2, according to various embodiments of the present invention; and
- (10) FIGS. 9A-9B set forth a flow diagram of method steps for encoding a source media title, according to various embodiments of the present invention.

DETAILED DESCRIPTION

- (11) In the following description, numerous specific details are set forth to provide a more thorough understanding of the various embodiments. However, it will be apparent to one of skilled in the art that the inventive concepts may be practiced without one or more of these specific details.
- (12) To optimize the overall visual experience that a media streaming service provides to viewers, the media streaming service provider oftentimes generates a bitrate ladder for each media title. The bitrate ladder for a given media title allows client-side endpoint devices to achieve a target visual quality during playback of the media title based on an available bandwidth. Each rung in the bitrate ladder specifies a different bitrate-resolution pair corresponding to a different pre-generated encoded version of the media title.
- (13) Some media streaming service providers use monolithic encoding techniques to generate the

different encoded versions of the media title, where each encoded version of the media title is associated with a different bitrate. Namely, the resolution and a rate control value used to encode the media content associated with a given media title do not vary across the media title. However, because the complexity of the media content associated with a given media title typically varies across the media title, the resulting encoded versions of the media title are often associated with encoding inefficiencies. More specifically, encoding relatively simple portions of the media title may consume more computational and storage resources than are necessary to meet a target visual quality. Further, the bandwidth required to stream the encoded versions of the media title may be unnecessarily large.

(14) For example, a movie could include relatively complex action portions and relatively simple monologues. The monologues could have the same visual quality regardless of whether the monologues are encoded using a bitrate of 3000 kilobits per second (kbps) or encoded using a bitrate of 560 kbps. By contrast, the required resources (e.g., computational resources, storage resources, bandwidth, etc.) associated with encoding the monologues using a bitrate of 3000 kbps exceed the required resources associated with encoding the monologues using a bitrate of 560 kbps. Accordingly, encoding the entire movie using a bitrate of 3000 kbps needlessly wastes computational and storage resources and unnecessarily increases the bandwidth required to stream the encoded version of the movie to endpoint devices.

(15) To reduce the encoding inefficiencies typically associated with monolithic encoding techniques, some media streaming service providers use a conventional subsequence-based encoding application to generate different encoded versions of a given media title. The conventional subsequence-based application varies the resolution and rate control value across the media title based on a target metric value, such as a target visual quality or a target bitrate. In this fashion, each encoded version of the media title is associated with a different target metric value.

(16) Typically, the conventional subsequence-based encoding application partitions the media title into different subsequences, where each subsequence is a shot that includes media content captured continuously from a given camera or other capture point. The conventional subsequence-based encoding application then encodes each subsequence numerous times at a variety of different encoding points to generate encoded subsequences. Each encoding point specifies a different combination of resolution and rate control value. Subsequently, the conventional subsequence-based encoding application performs optimization operations to generate different encoded versions of the media title.

(17) Finally, the conventional subsequence-based encoding application selects one of the encoded versions of the media title as the optimized encoded version of the media title based on the target metric value. The optimized encoded version of the media title includes a different combination of the encoded subsequences that span the length of the media title, and the resolution and rate control value may vary between the constituent encoded subsequences. As a result, the conventional subsequence-based encoding application reduces the encoding inefficiencies typically associated with monolithic encoding techniques described above.

(18) One drawback of this type of subsequence-based encoding technique is that the optimized encoded version of the media sequence does not necessarily comply with the transmission buffer constraints that are oftentimes associated with transmitting to and buffering at the endpoint devices. If an encoded version of a media sequence complies with the relevant transmission buffer constraints, then an associated endpoint device is going to be more able to properly buffer and play back the media title. If, however, the encoded version of the media sequence does not comply with the relevant transmission buffer constraints, then the endpoint device is going to be less able to properly buffer and play back media title in an uninterrupted fashion.

(19) One example of a common transmission buffer constraint is a constant bitrate (CBR) Video Buffer Verifier (VBV) constraint. A CBR VBV constraint specifies a transmission rate associated with a constant bandwidth channel and a buffer size associated with an endpoint device. Encoded

bits are streamed to an endpoint device via the constant bandwidth channel at the transmission rate. As encoded bits arrive at a decoder associated with the endpoint device, the encoded bits are stored in a transmission buffer having the buffer size. As part of displaying a particular frame of the media title, the decoder removes the associated encoded bits from the transmission buffer.

(20) Because the number of bits used to encode each frame can vary across the media title, the number of encoded bits that the transmission buffer stores can also vary during playback of the media title. If the optimized encoded media sequence is compliant with the CBR VBV constraint, then the variation in the number of bits used to encode each frame is limited and the endpoint device is going to be more able to properly buffer and play back the media title.

(21) If, however, the optimized encoded media sequence is not compliant with the CBR VBV constraint, then the variation in the number of bits used to encode each frame may cause the transmission buffer to underflow and/or overflow during playback of the media title. As a result, playback interruptions that adversely impact the associated viewing experience can occur. For example, at different points in time during the playback of the media title, the transmission buffer could be refilled faster than the transmission buffer was drained and, consequently, could overflow. If the transmission buffer were to overflow, then subsequent frames could be dropped, disrupting the playback of the media title.

(22) With the disclosed techniques, however, a media streaming service provider can perform subsequence-based encoding of a given media title subject to a transmission buffer constraint. Each encoded subsequence included in the optimized encoded media sequence is individually encoded at a potentially different optimized encoding point subject to the transmission buffer constraint. Consequently, the disclosed techniques can reduce encoding inefficiencies associated with monolithic encoding techniques while reducing the likelihood of playback interruptions relative to prior-art subsequence-based encoding techniques.

(23) In some embodiments, a buffer-based encoding application generates multiple encoded subsequences for each subsequence based on multiple encoding points. The buffer-based encoding application then incrementally generates encoded media sequences based on the encoded subsequences and the transmission buffer constraint in a playback order for the media title. More precisely, during an n .sup.th iteration, the buffer-based encoding application combines previously generated encoded media sequences spanning from the 0.sup.th subsequence to the $(n-1)$.sup.th subsequence with encoded subsequences associated with the n .sup.th subsequence. The resulting encoded media sequences span from the 0.sup.th subsequence to the n .sup.th subsequence.

(24) As the buffer-based encoding application generates each encoded media sequence, the buffer-based encoding application tracks the buffer occupancy of the transmission buffer. If the buffer-based encoding application determines that a variation in buffer occupancy for a particular encoded media sequence exceeds the size of the transmission buffer, then the buffer-based encoding application discards the encoded media sequence. The buffer-based encoding application also discards any encoded media sequences that are associated with a sub-optimal trade-off between average bitrate and an overall visual quality. After generating encoded media sequences that each span the length of the media title, the buffer-based encoding application selects the encoded media sequence having the highest overall visual quality as the optimized encoded media sequence.

(25) One advantage and technological advancement of the disclosed techniques relative to prior art techniques is that the optimized encoded media sequence is incrementally generated in a playback order based on sets of individually encoded subsequences and the transmission buffer constraint. Because each constituent encoded subsequence may be associated with a different encoding point, the overall visual quality level of the optimized encoded version of the media title is usually increased relative to a monolithically encoded version of the media title that is associated with the same average bitrate. Further, discarding encoded media sequences that do not comply with the transmission buffer constraint efficiently ensures that the optimized encoded media sequence complies with the transmission buffer constraint. As a result, the likelihood of degraded viewing

experiences attributable to transmission buffer limitations when streaming the media title to endpoint devices is reduced relative to prior art subsequence-based encoding techniques.

System Overview

(26) FIG. 1 is a conceptual illustration of a system **100** configured to implement one or more aspects of the present invention. As shown, the system **100** includes, without limitation, any number of compute instances **110** and a cloud **160**. In general, the cloud **160** contains encapsulated shared resources, software, data, etc. For explanatory purposes, multiple instances of like objects are denoted with reference numbers identifying the object and parenthetical numbers identifying the instance where needed.

(27) Any number of the compute instances **110** may reside outside the cloud **160** while other compute instances **110** may reside inside the cloud **160**. In various embodiments, any number of the components of the system **100** may be distributed across multiple geographic locations. Further, any number of the components of the system **100** may be included in or distributed across one or more stand-alone devices, distributed computing environments, or clouds **160** in any combination.

(28) As shown, each of the compute instances **110** includes, without limitation, a processor **112** and a memory **116**. The processor **112** may be any instruction execution system, apparatus, or device capable of executing instructions. For example, the processor **112** could comprise a central processing unit (CPU), a graphics processing unit (GPU), a controller, a microcontroller, a state machine, or any combination thereof. In alternative embodiments, each of the compute instances **110** may include any number of processors **112**.

(29) The memory **116** stores content, such as software applications and data, for use by the processor **112** of the compute instance **110**. The memory **116** may be one or more of a readily available memory, such as random access memory (RAM), read only memory (ROM), floppy disk, hard disk, or any other form of digital storage, local or remote. In some embodiments, a storage (not shown) may supplement or replace the memory **116**. The storage may include any number and type of external memories that are accessible to the processor **112**. For example, and without limitation, the storage may include a Secure Digital Card, an external Flash memory, a portable compact disc read-only memory (CD-ROM), an optical storage device, a magnetic storage device, or any suitable combination of the foregoing.

(30) In general, the compute instances **110** are configured to implement one or more applications. For explanatory purposes only, each application is depicted as residing in the memory **116** of a single compute instance **110** and executing on a processor **112** of the single compute instance **110**. However, as persons skilled in the art will recognize, the functionality of each application may be distributed across any number of other applications that reside in the memories **116** of any number of compute instances **110** and execute on the processors **112** of any number of compute instances **110** in any combination. Further, the functionality of any number of applications may be consolidated into a single application or subsystem.

(31) Although not shown, a media streaming service provides access to a library of media titles that can be viewed on a range of different endpoint devices. The library of media titles includes without limitation, the media title associated with the source media sequence **122**. Each endpoint device may connect to the media streaming service under different connection conditions that affect, without limitation, bandwidth and latency. In addition, each different endpoint device may include different hardware for outputting the media title to end user(s). For example, a given endpoint device could include a display screen having a particular screen size and a particular screen resolution.

(32) As described previously herein, in many embodiments, an endpoint device that connects to a media streaming service executes an endpoint application that determines, for a given media title, an appropriate encoded version of the media title to stream to the endpoint device based on the connection conditions and the properties of the endpoint device. More specifically, the endpoint application attempts to select a particular encoded version of the media title that is associated with

the best visual quality during playback of the media title on the endpoint device while avoiding playback interruptions due to buffering or re-buffering.

(33) The compute instances **110** are configured to generate encoded versions of media titles for streaming. More precisely, the compute instances **110** are configured to generate the optimized encoded media sequence **180** based on the source media sequence **122**. The source media sequence **122** includes, without limitation, any amount and type of media content that is associated with a media title. Examples of media content include, without limitation, any portion (including all) of feature length films, episodes of television programs, and music videos, to name a few. The optimized encoded media sequence **180** is a pre-generated encoded version of the media title and includes, without limitation, encoded media content derived from the media content included in the source media sequence **122**.

(34) As described in detail previously herein, the complexity and type of the media content associated with a given media title oftentimes varies across the media title. Using conventional monolithic encoding techniques, the resolution and the rate control value used to encode the media content do not vary across the media title. As referred to herein a “rate control value” is a value for a rate control parameter used to specify how an encoder is to allocate bits when performing encoding operations on media content in an associated rate control mode. One example of a rate control parameter is a quantization parameter (QP). As a result, encoding relatively simple portions of the media title may consume more computational and storage resources than what is necessary to meet the target visual quality. Among other things, such encoding inefficiencies waste computational and storage resources and increase the bandwidth required to stream encoded versions of media titles to endpoint devices.

(35) To reduce these types of encoding inefficiencies, conventional subsequence-based encoding techniques generate encoded versions of media titles in which the resolution and the rate control value vary across the media title. However, one limitation of conventional subsequence-based encoding techniques is that the encoded versions of a media title do not necessarily comply with the transmission buffer constraints that are oftentimes associated with transmitting to and buffering at the endpoint devices. As a result, playback interruptions, such as freeze frame and freeze drop, that disrupt the viewing experience can occur when those encoded media titles are streamed to endpoint devices for playback.

(36) More precisely, during streaming of an encoded version of the media title, as encoded bits arrive at a decoder associated with the endpoint device, the encoded bits are stored in a transmission buffer. As part of displaying a particular frame of the media title, the decoder removes the associated encoded bits from the transmission buffer. Because the number of bits used to encode each frame can vary across the media title, the number of encoded bits that the transmission buffer stores can also vary during playback of the media title. If the encoded version of the media title is compliant with the relevant associated transmission buffer constraints, then the variation in the number of bits used to encode each frame is limited and the endpoint device is going to be more able to properly buffer and play back the media title.

(37) If, however, the encoded version of a media title is not compliant with the relevant transmission buffer constraints, then the variation in the number of encoded bits used to encode each frame may cause the transmission buffer to overflow or underflow. Overflows and underflows of the transmission buffer can cause playback interruptions, such as freeze frame and freeze drop, that disrupt the viewing experience when those encoded media titles are streamed to endpoint devices for playback.

Encoding Individual Subsequences

(38) To comprehensively increase the quality of the viewing experience during streaming of a media title to an endpoint device relative to prior-art techniques, the system **100** includes, without limitation, a subsequence-based encoding subsystem **120**. The subsequence-based encoding subsystem **120** resides in the memory **116** of the compute instance **110(0)** and executes on the

processor **112** of the compute instance **110(0)**.

(39) The subsequence-based encoding subsystem **120** generates an optimized encoded media sequence **180** based on a transmission buffer constraint **170**. The transmission buffer constraint **170** includes, without limitation, a transmission rate **172**, a buffer size **174**, and a constraint type **176**. The transmission rate **172** specifies a rate at which encoded bits are transmitted to an endpoint device. The buffer size **174** specifies a minimum size of a transmission buffer associated with the endpoint device. The constraint type **176** specifies how encoded bits are transmitted and/or stored during streaming. In alternate embodiments, the system **100** may include any number and type of transmission buffer constraints **170** specified in any technically feasible fashion. In some such embodiments, the buffer-based encoding application **140** generates the encoded media sequence **180** based on all of the transmission buffer constraints **170**.

(40) One common constraint type **176** is a constant bit rate (CBR) Video Buffer Verifier (VBV) constraint. A CBR VBV constraint specifies a continuous transmission of encoded bits to a transmission buffer of a minimum buffer size **174** at the transmission rate **172**. Consequently, if the transmission buffer becomes full, then the endpoint device is unable to store some of the transmitted encoded bits. As persons skilled in the art will recognize, a CBR VBV constraint is usually a requirement for an encoded version of a media title to be streamed through a constant bandwidth channel, such as those used in broadcast television.

(41) Another common constraint type **176** is a variable bit rate (VBR) VBV constraint. A VBR VBV constraint specifies a conditional transmission of encoded bits to a transmission buffer of a minimum buffer size **174** at the transmission rate **172**. If the transmission buffer is not full, then the transmission of the encoded bits proceeds. If, however, the transmission buffer is full, then the transmission of the encoded bits pauses until the transmission buffer is no longer full. Some endpoint devices that include optical disk storage (e.g., Blu-Ray players) require that an encoded version of a media title is compliant with a VBR VBV constraint to ensure seamless playback of the media title.

(42) Notably, some hardware-based endpoint devices used in adaptive streaming require that encoded versions of media titles comply with VBV constraints because of limitations in the level of pre-buffering capabilities of the transmission buffers. In adaptive streaming, the endpoint device may switch between different encoded versions of a media title during playback of a media title. To facilitate adaptive streaming, any number of the techniques described herein may be executed multiple times based on different transmission buffer constraints **170** to generate multiple optimized encoded media sequences **160**. In alternative embodiments, the subsequence-based encoded subsystem **120** may be modified to generate a set of optimized encoded media sequences **180** based on a set of transmission buffer constraints **170**. In such embodiments, each optimized encoded media sequence **180(x)** is compliant with a different transmission buffer constraint **170(x)**.

(43) As shown, the subsequence-based encoding subsystem **120** includes, without limitation, a subsequence analyzer **130** and a buffer-based encoding application **140**. The subsequence analyzer **130** partitions the source media sequence **122** into any number of subsequences **132(0)-132(S)**. The subsequences **132** are non-overlapping, contiguous sets of frames that, together, span the source media sequence **122**. Each set of frames may represent a variety of different constructs, including a group of pictures (GOP), a sequence of frames, a plurality of sequences of frames, and so forth. For instance, in some embodiments, each subsequence **132** is a different shot included in the source media sequence **122**. As referred to herein, a “shot” is a sequence of frames that usually have similar spatial-temporal properties and run for an uninterrupted period of time.

(44) Upon receiving the source media sequence **122**, the subsequence analyzer **130** performs any number of analysis operations on the source media sequence **122** to determine the subsequences **132** based on any number and type of criteria. For instance, in various embodiments, the subsequence analyzer **130** may be configured to identify sets of frames for which a consistency metric lies within a specified range. In other embodiments, and as depicted in FIG. 1, the

subsequence analyzer **130** determines the subsequences **132** based on any number of shot changes **134**.

(45) More specifically, the subsequence analyzer **130** transmits the source media sequence **122** to a shot detector **124** that resides in the cloud **160**. The shot detector **124** executes any number of shot detection algorithms based on the source media sequence **122** to identify the shot changes **134**. Some examples of shot detection algorithms include, without limitation, a multi-scale sum-of-absolute-differences algorithm, a motion-compensated residual energy algorithm, a histogram of differences algorithm, a difference of histograms algorithm, and so forth. Each of the shot changes **134** specifies a boundary between a different pair of shots. The shot detector **124** then transmits the shot changes **134** to the subsequence analyzer **130**. Subsequently, the subsequence analyzer **130** performs partitioning operations on the source media sequence **122** based on the shot changes **134** to determine the subsequences **132**.

(46) In some embodiments, the subsequence analyzer **130** and/or the subsequence-based encoding subsystem **120** may perform any number type of additional operations as part of generating the subsequences **134**. For instance, in various embodiments, the subsequence analyzer **130** removes extraneous pixels from the source media sequence **122**. For example, the subsequence analyzer **130** could remove pixels included in black bars along border sections of the source media sequence **122**.

(47) In various embodiments, the subsequence-based encoding subsystem **120** ensures that the initial frame of each subsequence **132** is encoded as a key frame during encoding operations. As a general matter, a “key frame” and all subsequent frames from the same subsequence **132** that are included in an encoded media sequence are decoded independently of any proceeding frames included in the encoded media sequence.

(48) The subsequence-based encoding subsystem **120** may ensure that the different initial frames of the different subsequences **132** are encoded as key frames in any technically feasible fashion. For instance, in some embodiments, the buffer-based encoding application **140** configures an encoding application **162** to encode frames as key frames based on a key frame location list (not shown) when encoding media content. In other embodiments, the buffer-based encoding application **140** and/or the encoding application **162** may perform any number of encoding operations to encode the different initial frames of the different subsequences **132** as key frames when encoding media content.

(49) As persons skilled in the art will recognize, during playback, the media title associated with the source media sequence **122** is switchable between decoded versions of different optimized encoded media sequences **180** at aligned key frames to optimize a viewing experience based on any number of relevant criteria. Examples of relevant criteria include the current connection bandwidth, the current connection latency, the content of the upcoming subsequence **132**, and the like.

(50) As shown, the buffer-based encoding application **140** generates the optimized encoded media sequence **180** based on the subsequences **132** and the transmission buffer constraint **170**. The optimized encoded media sequence **180** includes, without limitation, S+1 encoded subsequences (not shown in FIG. 1) that are associated, respectively, with the subsequences **132(0)**-**132(S)**. Each of the encoded subsequences includes encoded media content derived from the media content included in the associated subsequence **132**.

(51) As described in greater detail in conjunction with FIG. 2, for each of the subsequences **132**, the buffer-based encoding application **140** generates multiple encoded subsequences based on an encoding point list. The encoding point list may include any number of encoding points. Each encoding point includes, without limitation, a resolution and a rate control value. The buffer-based encoding application **140** may generate the encoded subsequences in any technically feasible fashion. Further, as a general matter, the buffer-based encoding application **140** may generate encoded video content derived from video content based on a given resolution and given rate control value in any technically feasible fashion.

(52) For instance, in some embodiments, to generate the encoded subsequence corresponding to a given subsequence **132** and a given encoding point, the buffer-based encoding application **140** performs sampling operations on the subsequence based on the resolution to generate a sampled subsequence. The buffer-based encoding application **140** then configures the encoding application **162** to encode the sampled subsequence at the rate control value to generate the encoded subsequence.

(53) As shown, the encoding application **162** resides in the cloud **160** and is configured to efficiently perform encoding operations via one or more parallel encoders **164**. Each of the parallel encoders **164** may include any number of compute instances **110**. In alternative embodiments, the buffer-based encoding application **140** may perform encoding operations and the system **100** may omit the encoding application **162**. In the same or other embodiments, the system **100** may include a sampling application, and the buffer-based encoding application **140** may configure the sampling application to perform sampling operations.

(54) In various embodiments, the buffer-based encoding application **140** may perform sampling operations and encoding operations at any level of granularity (e.g., per frame, per subsequence **132**, per source media sequence **122**, etc.) in any combination and in any technically feasible fashion. For instance, in some embodiments, the buffer-based encoding application **140** may perform sampling operations on the source media sequence **122** based on a given resolution to generate a sampled media sequence. Subsequently, for each encoding point associated with the resolution, the buffer-based encoding application **140** may configure the encoding application **162** to encode the sampled media content corresponding to the associated subsequence **132** at the associated rate control value.

(55) For each encoded subsequence, the buffer-based encoding application **140** computes a bitrate, a quality score, and a distortion level. The buffer-based encoding application **140** may compute the bitrate, the quality score, and the distortion level for a given encoded subsequence in any technically feasible fashion. For each encoded subsequence, the buffer-based encoding application **140** then generates a subsequence encode point (not shown in FIG. 1). As described in greater detail in conjunction with FIG. 4, each subsequence encode point includes, without limitation, the encoded subsequence, the associated encoding point, the bitrate of the encoded subsequence, the quality score of the encoded subsequence, and the distortion level of the encoded subsequence.

(56) For each of the subsequences **132(x)**, the buffer-based encoding application **140** generates a different convex hull based on the subsequence encode points associated with the subsequence **132(x)**. In this fashion, the convex hull associated with a given subsequence **132(x)** includes any number of the subsequence encode points associated with the subsequence **132(x)**. In general, for a given subsequence **132**, the subsequence encode points included in the associated convex hull minimize the bitrate for different distortion levels.

(57) As described in greater detail in conjunction with FIGS. 2 and 7, the buffer-based encoding application **140** iteratively processes each of the convex hulls in a temporal playback order to generate a different associated global convex hull of media encode points. For explanatory purposes only, the “playback order” corresponds to the order in which an endpoint device plays back the subsequences **132**—from subsequence **132(0)** to subsequence **132(S)**.

(58) Each media encode point in the i.sup.th global convex hull is associated with a portion of the media title that spans from the 0.sup.th subsequence to the i.sup.th subsequence. In particular, each media encode point is associated with a different encoded media sequence. Each encoded media sequence includes, without limitation, (i+1) encoded subsequences, where each encoded subsequence is included in a different one of the convex hulls associated with the subsequences **132(0)-132(i)**. Importantly, each of the media encode points included in the i.sup.th global convex hull minimizes the average bitrate for a different overall distortion level subject to the transmission buffer constraint **170**.

(59) The buffer-based encoding application **140** determines the optimized encoded media sequence

180 based on the global convex hull associated with the subsequence **132(S)** that occurs last in the playback order. The global convex hull associated with the subsequence **132(S)** is also referred to herein as the “final” convex hull. The buffer-based encoding application **140** compares the average bitrates of the encoded media sequences associated with the final global convex hull and sets the optimized encoded media sequence **180** equal to the encoded media sequence having the highest average bitrate.

(60) As persons skilled in the art will recognize, the optimized encoded media sequence **180** has a higher overall quality level than any other encoded media sequence associated with the final global convex hull. Consequently, the optimized media sequence **180** is an encoded version of the media title that both complies with the transmission buffer constraint **170** and optimizes the overall visual quality during playback of the media title on an endpoint device. In alternative embodiments, the buffer-based encoding application **140** may select the optimized encoded media sequence **180** based on the final global convex hull and any number and type of criteria associated with any number and type of media metrics in any technically feasible fashion.

(61) As referred to herein, a media metric may be any measurement that corresponds to one or more properties of encoded video content, video content, audio content, and/or encoded audio content. Some examples of media metrics include, without limitation, bitrate, distortion metrics, audio quality metrics, visual quality metrics, etc. Example of visual quality metrics include, without limitation, a peak signal-to-noise ratio (PSNR), a linear video multimethod assessment fusion ((VMAF) metric, and a harmonic VMAF (VMAFh), to name a few.

(62) As described in greater detail in conjunction with FIG. 2, the buffer-based encoding application **140** optionally computes an optimized pre-buffering time **182** based on the optimized encoded media sequence **180** and the transmission buffer constraint **170**. In general, during streaming, a pre-buffering time specifies a total amount of time that an endpoint device waits after the streaming of a media title begins before beginning to play back a media title. During the pre-buffering time, the transmission buffer stores encoded bits and, consequently, when the endpoint device begins to play back the media title, the transmission buffer has an initial occupancy that is determined by the pre-buffering time.

(63) Finally, the buffer-based encoding application **140** transmits the optimized encoded media sequence **180** and, optionally, the optimized pre-buffering time **182** to the CDN **190** for distribution to endpoint devices. In alternative embodiments, the buffer-based encoding application **140** may cause the optimized pre-buffering time **182**, the optimized encoded media sequence **180**, and/or any number of the encoded subsequences to be delivered to endpoint devices in any technically feasible fashion. In the same or other embodiments, any amount and type of the functionality associated with the subsequence-based encoding application may be implemented in or distributed across any number of compute instances **110** and/or any number of endpoint devices, in any technically feasible fashion.

(64) For explanatory purposes only, the techniques described herein are described in the context of video encoding. However, as persons skilled in the art will recognize, the techniques described herein may be modified to optimize audio encoding instead of or in addition to video encoding. For instance, in some embodiments, an audio track may be partitioned into “audio scenes.” The audio scenes may be sampled via audio rendering hardware. The sampled audio scenes may be encoded via an audio encoder that is configured via a quantization parameter and/or bitrate setting. The quality scores of the encoded audio scenes may be computed via a perceptual audio quality metric, such as the Perceptual Evaluation of Audio Quality (PEAQ) algorithm. Notably, the audio encoder, any number of associated configuration parameters, the resolution and/or a rate control value may be optimized for each audio scene based on any of the techniques described herein in any combination.

(65) Further, although the techniques described herein are described in the context of media streaming, the techniques described herein may be modified to optimize encoding for any type of

media consumption. For example, the techniques may be altered to optimize encoding for online video gaming, screen-content sharing, two-way video conferencing, electronic communication, etc. (66) Note that the techniques described herein are illustrative rather than restrictive, and may be altered without departing from the broader spirit and scope of the invention. Many modifications and variations on the functionality provided by the subsequence-based encoding subsystem **120**, the subsequence analyzer **130**, the buffer-based encoding application **140**, the shot detector **124**, the encoding application **162**, and the content delivery network **190** will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the described embodiments. As a general matter, the techniques outlined herein are applicable to generating an encoded media sequence based on individually encoded subsequences and any number and type of transmission buffer constraints in any technically feasible fashion.

(67) It will be appreciated that the system **100** shown herein is illustrative and that variations and modifications are possible. For example the functionality provided by the subsequence-based encoding subsystem **120**, the subsequence analyzer **130**, the buffer-based encoding application **140**, the shot detector **124**, the encoding application **162**, and the content delivery network **190** as described herein may be integrated into or distributed across any number of software applications (including one), hardware devices (e.g., a hardware-based encoder), and any number of components of the system **100**. Further, the connection topology between the various units in FIG. **1** may be modified as desired.

(68) FIG. **2** is a more detailed illustration of the buffer-based encoding application **140** of FIG. **1**, according to various embodiments of the present invention. As shown, the buffer-based encoding application **140** includes, without limitation, an encoding point list **210**, subsequence encode sets **220**, a convex hull generator **230**, convex hulls **240**, a subsequence iterator **250**, global convex hulls **260**, and a selection engine **290**. The number of the subsequence encode sets **220** equals the number of subsequences **132**. Similarly, the number of the convex hulls **240** equals the number of subsequences **132**. In general, the subsequence **132(x)** is associated with the subsequence encode set **220(x)** and the convex hull **240(x)**.

(69) As shown, the encoding point list **210** includes any number of encoding points **212(0)-212(P)**. As described in greater detail in conjunction with FIG. **3**, each of the encoding points **212** includes, without limitation, a resolution and a rate control value. In alternative embodiments, each of the encoding points **212** may specify any number and type of parameters that impact encoding operations in any technically feasible fashion. For instance, in some embodiments, each of the encoding points **212** may specify a value for a “speed” configuration parameter associated with the encoding application **162**.

(70) In operation, the buffer-based encoding application **140** generates the subsequence encode sets **220(0)-220(S)** based on the subsequences **132(0)-132(S)** and the encoding point list **210**. Each of the subsequence encode sets **220(0)-220(S)** includes, without limitation, subsequence encode points **222(0)-222(P)**. The number of subsequence encode points **222** included in each of the subsequence encode sets **220** equals the number of the encoding points **212** included in the encoding point list **210**. As described in greater detail in conjunction with FIG. **4**, each of the subsequence encode points **222(x)** includes, without limitation, the associated encoding point **212(x)**, an encoded subsequence, a bitrate, a quality score, and a distortion level.

(71) For each combination of subsequence **132(x)** and encoding point **212(y)**, the buffer-based encoding application **140** generates a corresponding subsequence encode point **222(y)** and adds the subsequence encode point **222(y)** to the subsequence encode set **220(x)**. More specifically, the buffer-based encoding application **140** configures the encoding application **162** to encode the subsequence **132(x)** at the resolution and rate control value included in the encoding point **212(y)**. The buffer-based encoding application **140** then computes a bitrate of the encoded subsequence, a quality score of the encoded subsequence, and a distortion level of the encoded subsequence. As a result, the buffer-based encoding application **140** generates $(S+1) \times (P+1)$ different subsequence

encode points **222**.

(72) In alternative embodiments, each of the subsequences **132(x)** may be associated with a different encoding point list **210(x)** and the number of encoding points **212** in the encoding point list **210(x)** may differ from the number of encoding points **212** in any of the other encoding point lists **210**. In a complementary fashion, the number of subsequence encode points **222** included in the subsequence encode set **220(x)** may differ from the number of subsequence encode points **222** included in any of the other subsequence encode sets **220**.

(73) The buffer-based encoding application **140** may generate each encoded subsequence and determine the associated bitrate, the associated quality score, and the associated distortion level in any technically feasible fashion. For instance, in some embodiments, to generate the encoded subsequence associated with both the subsequence **132(x)** and the encoding point **212(y)**, the buffer-based encoding application **140** first performs sampling operations to generate a sampled subsequence. More specifically, the buffer-based encoding application **140** samples the subsequence **132(x)** based on the resolution included in the encoding point **212(y)** to generate a sampled subsequence. Subsequently, the buffer-based encoding application **140** causes the encoding application **162** to encode the sampled subsequence using the rate control value included in the encoding point **212(y)** to generate the encoded subsequence.

(74) In some embodiments, to determine the quality score of the encoded subsequence, the buffer-based encoding application **140** decodes the encoded subsequence to generate a decoded subsequence. The buffer-based encoding application **140** then re-samples (i.e., up-samples or down-samples) the decoded subsequence to a target resolution to generate a re-constructed subsequence that is relevant to the display characteristics of a class of endpoint devices.

(75) In alternative embodiments, the buffer-based encoding application **140** may compute any number of quality scores for any number of resolutions. For example, a certain video may be delivered in 3840×2160 resolution, yet be intended to be consumed by a large number of displays in 1920×1080 resolution. Another class of endpoint devices, for example laptop computers, is expected to display the same video in 1280×720 resolution. Yet another class of endpoint devices, for example, tablet or smartphone devices, is expected to display the same video in 960×540 resolution. The buffer-based encoding application **140** could up-sample the decoded subsequence to all these target resolutions in order to assess quality, when considering one of these different classes of endpoint devices, correspondingly.

(76) The buffer-based encoding application **140** then analyzes the reconstructed subsequence to generate the quality score for a quality metric (QM). For instance, in some embodiments the buffer-based encoding application **140** implements a VMAF (or harmonic VMAF) algorithm to generate a VMAF score for each encoded subsequence based on the associated re-constructed subsequence. Although a multitude of video quality metrics, such as VMAF scores, can be calculated at different target resolutions, it should be clear that, when comparing qualities among encoded subsequences associated with different resolutions, applications need to use the same target resolution for re-sampling, after decoding. For instance, in some embodiments, the buffer-based encoding application **140** re-samples the decoded subsequence to 1920×1080 to generate a re-constructed subsequence. Subsequently, the buffer-based encoding application **140** computes the quality score for the encoded subsequence based on the associated re-constructed subsequence.

(77) The buffer-based encoding application **140** may generate the bitrate in any technically feasible fashion. For instance, in some embodiments, the buffer-based encoding application **140** may divide the total number of encoded bits included in the encoded subsequence by the length of the associated subsequence **132**. In the same or other embodiments, the buffer-based encoding application **140** may compute the distortion level based on the quality score and any technically feasible technique for converting quality scores to distortion levels. For example, the buffer-based encoding application **140** could invert the quality score to determine the distortion level. In another example, the buffer-based encoding application **140** could subtract the quality score from a

constant value to determine the distortion level.

(78) For each of the subsequences **132(x)**, the convex hull generator **230** generates a convex hull **240(x)** based on the subsequence encode set **220(x)**. Each of the convex hulls **240(x)** includes, without limitation, the subsequence encode points **222** included in the subsequence encode set **220(x)** that minimize bitrate for a given distortion level. A detailed example of how the convex hull generator **230** generates the convex hull **240(0)** based on the subsequence encode set **220(0)** is described in conjunction with FIG. 5.

(79) The subsequence iterator **250** processes the convex hulls **240** sequentially, in a playback order for the media title. More specifically, for an index i that starts at 0 and ends at S , an instance of the subsequence iterator **250(i)** processes the convex hull **240(i)** associated with the subsequence **132(i)**. As described in greater detail in conjunction with FIG. 7, the subsequence iterator **250(i)** generates the global convex hull **260(i)** based on the global convex hull **260(i-1)** and the convex hull **240(i)**. The global convex hull **260(i)** is associated with the subsequences **132(0)-132(i)** and includes, without limitation, any number of media encode points **270**. The number of media encode points **270** included in the global convex hull **260(i)** may differ from the number of media encode points **270** in any of the other global convex hulls **260**.

(80) As described in greater detail in conjunction with FIG. 6, each of the media encode points **270** includes, without limitation, a subsequence encode list, an encoded media sequence, an average bitrate, an overall distortion level, a current buffer occupancy, a minimum buffer occupancy, and a maximum buffer occupancy. For each of the media encode points **270** included in the global convex hull **260(i)**, the subsequence encode list specifies $(i+1)$ subsequence encode points. More precisely, the subsequence encode list specifies a different subsequence encode point for each of the subsequences **132(0)-132(i)**. The encoded media sequence includes, without limitation, the $(i+1)$ encoded subsequences included in the $(i+1)$ subsequence encode points specified in the subsequence encode list. The average bitrate and the overall distortion level specify, respectively, a bitrate of the encoded media sequence and a distortion level of the encoded media sequence.

(81) The current buffer occupancy, the minimum buffer occupancy, and the maximum buffer occupancy specify characteristics of an exemplary transmission buffer when streaming the encoded media sequence to an exemplary endpoint device as per the transmission buffer constraint **170**. The current buffer occupancy specifies the number of bits that are stored in the exemplary transmission buffer when the encoded media sequence is decoded for playback on the exemplary endpoint device. The minimum buffer occupancy and the maximum buffer occupancy specify, respectively, the minimum and maximum number of bits that are stored in the exemplary transmission buffer during streaming of the encoded media sequence to the exemplary endpoint device.

(82) Although not shown, to generate the global convex hull **260(0)**, the buffer-based encoding application **140** initializes the subsequence iterator **250(0)**. More precisely, the buffer-based encoding application **140** generates an initial global convex hull **260(-1)** that includes a single initialization media encode point **270**. The buffer-based encoding application **140** sets the subsequence encode list included in the initialization media encode point **270** equal to an empty set and the encoded subsequence included in the initialization media encode point **270** equal to NULL. The buffer-based encoding application **140** sets each of the current buffer occupancy, the minimum buffer occupancy, and the maximum buffer occupancy included in the initialization media encode point **270** equal to the buffer size **174**. In alternative embodiments, the buffer-based encoding application **140** may initialize the subsequence iterator **250(0)** in any technically feasible fashion. For instance, in some embodiments, the buffer-based encoding application **140** may initialize the subsequence iterator **250(0)** based on a specified pre-buffering time.

(83) In operation, to generate the global convex hull **260(i)**, the subsequence iterator **250(i)** combines each of the subsequence encode lists associated with the global convex hull **260(i-1)** with each of the subsequence encode points included in the convex hull **240(i)** to generate new subsequence encode lists. Accordingly, if the global convex hull **260(i-1)** includes N media encode

points **270** and the convex hull **240(i)** includes M subsequence encode points, then the subsequence iterator **250(i)** generates $(N \times M)$ new subsequence encode lists.

(84) For each new subsequence encode list, the subsequence iterator **250(i)** generates a new media encode point **270**. More precisely, for a given new subsequence encode list, the subsequence iterator **250(i)** aggregates the different encoded subsequences included in the subsequence encode points specified in the subsequence encode list to generate an encoded media sequence. The subsequence iterator **250(i)** computes the average bitrate and the overall distortion level based on the encoded media sequence. The subsequence iterator **250(i)** then computes the current buffer occupancy, the minimum buffer occupancy, and the maximum buffer occupancy based on encoded media sequence and the transmission buffer constraint **170**. Subsequently, the subsequence iterator **250(i)** generates a new media encode point **270** that includes the new subsequence encode list, the associated encoded media sequence, the average bitrate, the overall distortion level, the current buffer occupancy, the minimum buffer occupancy, and the maximum buffer occupancy. The subsequence iterator **250(i)** adds the new media encode point **270** to a media encode set (not shown in FIG. 2).

(85) Subsequently, the subsequence iterator **250(i)** performs filtering operations on the media encode set based on the transmission buffer constraint **170** to generate a compliant media encode set (not shown in FIG. 1). As part of the filtering operations, the subsequence iterator **250(i)** discards any media encode points **270** included in the media encode set that do not comply with the transmission buffer constraint **170**. The subsequence iterator **250(i)** then discards any media encode points **270** included in the compliant media encode set that are associated with a sub-optimal trade-off between average bitrate and an overall visual quality to generate the global convex hull **260(i)**.

(86) Importantly, for the portion of the media title that includes the subsequences **132(0)-132(i)**, each of the media encode points **270** included in the global convex hull **260(i)** minimizes the average bitrate for a different overall distortion level subject to the transmission buffer constraint **170**. Further, because the global convex hull **260(i)** typically includes significantly fewer media encode points than the associated media encode set, the number of subsequence encode points included in each of the subsequent subsequence media encode sets is reduced. Consequently, the time required to generate the subsequent global convex hulls **260** is optimized.

(87) For example, suppose that the media encode set were to include **100** media encode points **270** and the subsequence iterator **250(i)** were to discard **60** of the media encode points **270** based on the transmission buffer constraint **170**. Further, suppose that the subsequence iterator **250(i)** were to determine that 30 of the 40 remaining media encode points **270** were associated with sub-optimal trade-offs between an average bitrate and an overall visual quality. The resulting global convex hull **260(i)** would include 10 media encode points **270**. Subsequently, the subsequence iterator **250(i+1)** would combine the 10 media encode points **270** included in the global convex hull **260(i)** with A points included in the convex hull **240(i+1)** to generate a new media encode set that would include $(10 \times A)$ media encode points **270**.

(88) As shown, the selection engine **290** determines the optimized encoded media sequence **180** based on the global convex hull **260(S)**. As described previously herein, the global convex hull **260(S)** is associated with all the subsequences **132(0)-132(S)**. To determine the optimized media sequence **180**, the selection engine **290** compares the average bitrates of the media encode points **270** included in the global convex hull **260(S)** to determine an optimized media encode point **270** having the highest average bitrate. The selection engine **290** then sets the optimized encoded media sequence **180** equal to the encoded media sequence included in the optimized media encode point **270**.

(89) As persons skilled in the art will recognize, the optimized encoded media sequence **180** has a higher overall quality level than any other encoded media sequence associated with the global convex hull **260(S)**. Consequently, the optimized media sequence **180** is an encoded version of the media title that both complies with the transmission buffer constraint **170** and optimizes the overall

visual quality during playback of the media title on an endpoint device. In alternative embodiments, the selection engine **290** may select the optimized encoded media sequence **180** based on the global convex hull **260(S)** and any number and type of criteria associated with any number and type of media metrics in any technically feasible fashion.

(90) In some embodiments, the subsequence iterator **250** does not generate the encoded media sequences included in the media encode points **270**. For each media encode point **270**, the subsequence iterator **250** determines properties of the associated encoded media subsequence based on the associated subsequence encode list. In particular, the subsequence iterator **250** computes the average bitrate, the overall distortion level, the current buffer occupancy, the minimum buffer occupancy, and the maximum buffer occupancy based on the subsequence encode list. After selecting the optimized media encode point **270**, the selection engine **290** generates the optimized encoded media sequence **180** based on the subsequence encode list included in the optimized media encode point **270**. More precisely, the selection engine **290** aggregates the different encoded subsequences specified in the subsequence encode list included in the optimized media encode point **270** to generate the optimized encoded media sequence **180**. In this fashion, each subsequence encode list is an “encoding recipe” that specifies how the associated encoded media subsequence is generated.

(91) After the selection engine **290** determines the optimized encoded media sequence **180**, the selection engine **290** determines whether the transmission buffer constraint **170** is associated with a configurable pre-buffering time. In general, a pre-buffering time specifies a total number of encoded bits that an endpoint device is to store before beginning to play back a media title. The selection engine **290** may determine whether the transmission buffer constraint **170** is associated with a configurable pre-buffering time in any technically feasible fashion.

(92) For instance, in some embodiments, if the constraint type **176** is CRB VBV, then the selection engine **290** determines that the transmission buffer constraint **170** is associated with a configurable pre-buffering time. Otherwise, the selection engine **290** determines that the transmission buffer constraint **170** is not associated with a configurable pre-buffering time. In other embodiments, the selection engine **290** determines whether the transmission buffer constraint **170** is associated with a configurable pre-buffering time based on user input received via a graphical user interface (not shown).

(93) If the selection engine **290** determines that the transmission buffer constraint **170** is associated with a configurable pre-buffering time, then the selection engine **290** computes an optimized pre-buffering time **182** based on the optimized media encode point **270**. More precisely, the selection engine **290** sets an initial buffer occupancy equal to the difference between the buffer size **174** and the minimum buffer occupancy included in the optimized media encode point **270**. The selection engine **290** then sets the optimized pre-buffering time **182** equal to the quotient of the initial buffer occupancy and the transmission rate **172**. In alternative embodiments, the buffer-based encoding application **140** may compute the optimized pre-buffering time **182** and/or an initial buffer occupancy in any technically feasible fashion. The selection engine **290** then transmits the optimized pre-buffering time **182** to a content delivery network (CDN) **190** for distribution to endpoint devices in conjunction with the optimized encoded media sequence **180**.

(94) Irrespective of whether the transmission buffer constraint **170** is associated with a configurable pre-buffering time, the selection engine **290** transmits the optimized encoded media sequence **180** to the CDN **190** for distribution to endpoint devices. In various embodiments, the selection engine **290** may be configured to identify each of the individual encoded subsequences that are included in the optimized encoded media sequences **180**. The selection engine **290** may then transmit the identified encoded subsequences to the CDN **190** for distribution to endpoint devices.

(95) In alternative embodiments, the selection engine **290** may compute and transmit any number and type of encoded subsequences, configuration parameters, and the like to the CDN **190** or any number of endpoint devices in any technically feasible fashion. For instance, in some

embodiments, the selection engine **190** may transmit an initial buffer occupancy to the CDN **190** instead of the optimized pre-buffering time **182**. In other embodiments, the selection engine **290** may transmit the optimized encoded media sequence **180** and, optionally, the optimized pre-buffering time **182** to any number of endpoint devices in any technically feasible fashion. For example, the selection engine **290** could signal the pre-buffering time through a system header as part of streaming the optimized encoded media sequence **180** to endpoint devices.

(96) FIG. **3** is a more detailed illustration of the encoding point list **210** of FIG. **2**, according to various embodiments of the present invention. As shown, the encoding point list **210** includes any number of encoding points **212(0)-212(P)**. Each of the encoding points **212** includes, without limitation, a resolution **314** and a rate control value **316**. The rate control value **316** may be a value for any parameter that specifies a tradeoff between bitrate and distortion level or quality during encoding. For instance, in some embodiments, the rate control value **316** is a value for a quantization parameter (QP) that allows a monotonic performance in terms of bitrate and distortion level when encoding video content. The higher the “QP,” the lower the resulting bitrate at the expense of lower quality.

(97) For explanatory purposes only, resolutions **314**, and rate control values **316** for encoding points **212(0)-212(3)**, **212(P-1)**, and **212(P)** are depicted in italics. As shown, the encoding point **212(0)** includes the resolution **314** of 640×360 and the rate control value **316** QP=51. The encoding point **212(1)** includes the resolution **314** of 960×540 and the rate control value **316** QP=51. The encoding point **212(2)** includes the resolution **314** of 1280×720 and the rate control value **316** QP=51. The encoding point **212(3)** includes the resolution **314** of 1920×1080 and the rate control value **316** QP=51. The encoding point **212(P-1)** includes the resolution **314** of 1280×720 and the rate control value **316** QP=0. The encoding point **212(P)** includes the resolution **314** of 1920×1080 and the rate control value **316** QP=0.

(98) As the depicted exemplary encoding points **212** illustrate, any combination of the resolution **314** and the rate control value **316** may differ between any two encoding points **212**. In general, the buffer-based encoding application **140** may acquire or generate the encoding list **210** in any technically feasible fashion. For instance, in some embodiments, the buffer-based encoding application **140** may generate the encoding list **210** based on permutations of eight resolutions **314** and all the rate control values **316** allowed by the encoding application **162**. Advantageously, because of the wide variety of encoding points **212**, the buffer-based encoding application **140** generates a wide variety of different encoded subsequences for each subsequence **132**.

(99) FIG. **4** is a more detailed illustration of one of the subsequence encode sets **220** of FIG. **2**, according to various embodiments of the present invention. The subsequence encode set **220(0)** is associated with the subsequence **132(0)**. As shown, the subsequence encode set **220(0)** includes, without limitation, subsequence encode points **222(0)-222(P)**. In general, the number of subsequence encode points **222** included in the subsequence encode set **220** is equal to the number of encoding points **212** included in the encoding point list **210**. Further, the subsequence encode point **222(x)** is associated with the encoding point **212(x)**.

(100) Each subsequence encode point **212** includes, without limitation, the associated encoding point **212**, an encoded subsequence **430**, a bitrate **432** of the encoded subsequence **430**, a quality score **434** of the encoded subsequence **430**, and a distortion level **436** of the encoded subsequence **430**. As described previously herein in conjunction with FIG. **2**, the buffer-based encoding application **140** may generate the encoded subsequence **430** in any technically feasible fashion based on the associated encoding point **212** and the associated subsequence **132**. Subsequently, the buffer-based encoding application **140** may determine the bitrate **432**, the quality score **434**, and the distortion level **436** in any technically feasible fashion based on the encoded subsequence **430**. The quality score **434** may be the value for any quality metric.

(101) For explanatory purposes only, exemplary values for the encoding point **212(0)** included in the subsequence encode point **222(0)** and the encoding point **212(P)** included in the subsequence

encode point **222(P)** are depicted in italics. The encoding point **212(0)** includes, without limitation, the resolution **314** of 640×360, and the rate control value **316** QP=51. Accordingly, the buffer-based encoding application **140** configures the encoding application **162** to encode the subsequence **132(0)** at the resolution **314** of 640×360 and the rate control value **316** QP=51 to generate the encoded subsequence **430(0)**. By contrast, the encoding point **212(P)** includes, without limitation, the resolution **314** of 1920×1080 and the rate control value **316** QP=0. Accordingly, the buffer-based encoding application **140** configures the encoding application **162** to encode the subsequence **132(0)** at the resolution **314** of 1920×1080, and the rate control value **316** QP=0 to generate the encoded subsequence **430(P)**.

(102) FIG. 5 illustrates an exemplary convex hull **240(0)** that is generated by the convex hull generator **230** of FIG. 2, according to various embodiments of the present invention. In particular, the convex hull generator **230** generates the convex hull **240(0)** based on the subsequence encode set **220(0)**. As shown, a graph **500** includes, without limitation, a bitrate axis **510** and a distortion axis **520**.

(103) In operation, for each of the subsequence encode points **222** included in the subsequence encode set **220(0)** the convex hull generator **230** generates a corresponding plotted subsequence encode point **222** in the graph **500**. The convex hull generator **230** plots a given subsequence encode point **222** by locating the bitrate **432** along the bitrate axis **510** and the distortion level **436** along the distortion axis **520**. The convex hull generator **230** then evaluates the plotted subsequence encode points **222** to determine the convex hull **240(0)**.

(104) More specifically, the convex hull generator **230** identifies the plotted subsequence encode points **222** that form a boundary where all the plotted subsequence encode points **222** reside on one side of the boundary (in this case, the right side of the boundary) and also are such that connecting any two consecutive identified plotted subsequence encode points **222** with a straight line leaves all remaining plotted subsequence encode points **222** on the same side. The convex hull **240(0)** includes the set of the identified subsequence encode points **222**.

(105) For explanatory purposes only, the subsequence encode points **222** that are included in the convex hull **240(0)** are depicted as crosses in the graph **500**, while the remaining subsequence encode points **222** are depicted as filled circles in the graph **500**. In particular, the subsequence encode point **222(P)** is included in the convex hull **240(0)**, while the subsequence encode point **222(0)** is not included in the convex hull **240(0)**.

(106) Persons skilled in the art will understand that many techniques for generating convex hulls are well known in the field of mathematics, and all such techniques may be implemented to generate the convex hulls **240**. In one embodiment, the convex hull generator **230** applies machine-learning techniques to estimate the subsequence encode points **222** included in the convex hull **240** based on various parameters of the associated subsequence **132** and/or source media sequence **122**. In this manner, some of the computations discussed thus far may be streamlined and/or avoided entirely.

Combining Different Encoded Subsequences

(107) FIG. 6 is a more detailed illustration of one of the media encode points **270** of FIG. 2, according to various embodiments of the present invention. For explanatory purposes only, the media encode point **270(k)** depicted in FIG. 6 is included in the global convex hull **260(i)** that is associated with the subsequences **132(0)-132(i)**. As shown, the media encode point **270(k)** includes, without limitation, a subsequence encode list **610**, an encoded media sequence **630**, an average bitrate **632**, an overall distortion level **636**, a current buffer occupancy **650**, a minimum buffer occupancy **652**, and a maximum buffer occupancy **654**.

(108) The subsequence encode list **610** includes, without limitation, encode selections **620(0)-620(i)**, where the number of encode selections **620** is equal to the number of subsequences **132** associated with the global convex hull **260(i)**. More specifically, the subsequence encode list **610** includes a different encode selection **620(x)** for each subsequence **132(x)** associated with the

global convex hull **260(i)**. Each of the encode selections **620** specifies a different subsequence encode point **222**. In general, the encode selection **620(x)** specifies one of the subsequence encode points **222** included in the subsequence encode set **220(x)** associated with the subsequence **132(x)**. (109) For explanatory purposes only, exemplary values for the encode selections **620(0)**, **620(1)**, and **620(i)** are depicted in italics. The encode selection **620(0)** specifies the subsequence encode point **222(5)** in the subsequence encode set **220(0)**. The encode selection **620(1)** specifies the subsequence encode point **222(0)** in the subsequence encode set **220(1)**. The encode selection **620(i)** specifies the subsequence encode point **222(P)** in the subsequence encode set **220(i)**.

(110) The encoded media sequence **630** includes, without limitation, the (i+1) encoded subsequences **430** included in the (i+1) subsequence encode points **222** specified in the subsequence encode list **610**. The average bitrate **632** and the overall distortion level **636** specify, respectively, a bitrate of the encoded media sequence **630** and a distortion level of the encoded media sequence **630**. The subsequence iterator **250(i)** and/or the buffer-based encoding application **140** may determine the encoded media sequence **630**, the average bitrate **632**, and the overall distortion level **636** in any technically feasible fashion.

(111) As described in conjunction with FIG. 2, the current buffer occupancy **650**, the minimum buffer occupancy **652**, and the maximum buffer occupancy **654** specify characteristics of an exemplary transmission buffer having a size equal to the buffer size **174**. More precisely, the current buffer occupancy **650** specifies the number of bits that are stored in the exemplary transmission buffer when the encoded media sequence **630** is decoded for playback on the exemplary endpoint device. The minimum buffer occupancy **652** specifies the minimum of bits that are stored in the exemplary transmission buffer during streaming of the encoded media sequence **630** to the exemplary endpoint device. The maximum buffer occupancy **654** specifies the maximum of bits that are stored in the exemplary transmission buffer during streaming of the encoded media sequence **630** to the exemplary endpoint device.

(112) FIG. 7 is a more detailed illustration of the subsequence iterator **250(i)** of FIG. 2, according to various embodiments of the present invention. In operation, the subsequence iterator **250(i)** generates the global convex hull **260(i)** based on the global convex hull **260(i-1)**, the convex hull **240(i)** and the transmission buffer constraint **170**. For explanatory purposes only, the global convex hull **260(i-1)** includes, without limitation V media encode points **270**. As described in conjunction with FIG. 6, each of the media encode points **270** includes a different subsequence encode list **610**. Accordingly, the global convex hull **260(i-1)** is associated with V subsequence encode lists **610**. The convex hull **240(i)** includes, without limitation, A subsequence encode points **222**. As shown, the subsequence iterator **250(i)** includes, without limitation, a combination engine **710**, a filtering engine **730**, and a global convex hull generator **750**.

(113) The combination engine **710** generates a media encode set **720** based on the global convex hull **260(i-1)**, the convex hull **240(i)** and the transmission buffer constraint **170**. First, the combination engine **710** combines each of the V subsequence encode lists **610** associated with the global convex hull **260(i-1)** with each of the A subsequence encode points **222** included in the convex hull **240(i)** to generate (A×V) new subsequence encode lists **610**. To combine the subsequence encode list **610(x)** with the subsequence encode point **222(y)**, the combination engine **710** appends the subsequence encode point **222(y)** to the subsequence encode list **610(x)** to generate a new subsequence encode list **610(z)**. Notably, each of the subsequence encode lists **610** associated with the global convex hull **260(i-1)** is associated with the subsequences **132(0)-132(i-1)** and each of the subsequence encode points **222** included in the convex hull **240(i)** is associated with the subsequence **132(i)**. Consequently, each of the new subsequence encode lists **610** is associated with the subsequences **132(0)-132(i+1)**.

(114) For each new subsequence encode list **610**, the subsequence iterator **250(i)** generates a new media encode point **270** based on the subsequence encode list **610** and then adds the new media encode point **270** to the media encode set **720**. More precisely, for the new subsequence encode list

610(z), the subsequence iterator **250(i)** aggregates the different encoded subsequences **430** included in the subsequence encode points **222** specified in the subsequence encode list **610(z)** to generate the encoded media sequence **630(z)**. Subsequently, the subsequence iterator **250(i)** computes the average bitrate **632(z)** and the overall distortion level **636(z)** based on the encoded media sequence **630(z)**. The subsequence iterator **250(i)** may compute the average bitrate **632(z)** and the overall distortion level **636(z)** in any technically feasible fashion.

(115) For instance, in some embodiments, the subsequence iterator **250(i)** weights the bitrate **432** of each of the encoded subsequences **430** included in the encoded media sequence **630(z)** by the associated duration to generate weighted bitrates. The subsequence iterator **250(i)** sums the weighted bitrates and then divides the resulting weighted sum by the total duration of the encoded media sequence **630(z)** to compute the average bitrate **632(z)**. The subsequence iterator **250(i)** computes the overall distortion level **636(z)** in a similar fashion.

(116) In alternative embodiments, the subsequence iterator **250(i)** computes the average bitrate **632(z)** based on the subsequence encode list **610(x)** and the subsequence encode point **222(y)** from which the subsequence encode list **610(z)** was derived. More precisely, the subsequence iterator **250(i)** weights the average bitrate **632(x)** associated with the subsequence encode list **610(x)** by the duration of the encoded media sequence **630(x)** to generate a weighted media bitrate. Similarly, the subsequence iterator **250(i)** weights the bitrate **432** included in the subsequence encode point **222(y)** by the duration of the encoded subsequence **430(y)** to generate a weighted subsequence bitrate. The subsequence iterator **250(i)** sums the weighted media bitrate and the weighted subsequence bitrate and then divides the resulting weighted sum by the total duration of the encoded media sequence **630(z)** to compute the average bitrate **632(x)**. The subsequence iterator **250(i)** computes the overall distortion level **636(z)** in a similar fashion.

(117) The subsequence iterator **250(i)** computes the current buffer occupancy **650(z)**, the minimum buffer occupancy **652(z)**, and the maximum buffer occupancy **654(z)** based on encoded media sequence **630(z)** and the transmission buffer constraint **170**. The current buffer occupancy **650(z)** specifies the number of bits that are stored in an exemplary transmission buffer when the encoded media sequence **630(z)** is decoded for playback on the exemplary endpoint device. The minimum buffer occupancy **652(z)** and the maximum buffer occupancy **654(z)** specify, respectively, the minimum and maximum number of bits that are stored in the exemplary transmission buffer during streaming of the encoded media sequence **630(z)** to the exemplary endpoint device as per the transmission buffer constraint **170**. The subsequence iterator **250(i)** may compute the current buffer occupancy **650(z)**, the minimum buffer occupancy **652(z)**, and the maximum buffer occupancy **654(z)** in any technically feasible fashion that is consistent with the transmission buffer constraint **170**.

(118) For instance, in some embodiments, the subsequence iterator **250(i)** first computes the current buffer occupancy **650(z)** based on the subsequence encode list **610(x)** and the subsequence encode point **222(y)** from which the new subsequence encode list **610(z)** was derived. More precisely, the subsequence iterator **250(i)** sets a number of drained bits equal to the number of encoded bits included in the encoded subsequence **430(y)** included in the subsequence encode point **222(y)**. The subsequence iterator **250(i)** also sets a number of refilled bits count equal to the product of the transmission rate **172** and the playback duration of the encoded subsequence **430(y)**. The subsequence iterator **250(i)** then subtracts the number of drained bits from the number of refilled bits to determine a subsequence delta.

(119) The subsequence iterator **250(i)** sets a potential buffer occupancy equal to the sum of the subsequence delta and the current buffer occupancy **650(x)** associated with the subsequence encode list **610(x)**. If the constraint type **176** is CBR VBV, then the subsequence iterator **250(i)** sets the current buffer occupancy **650(z)** equal to the potential buffer occupancy. If, however, the constraint type **176** is VBR VBV, then the subsequence iterator **250(i)** sets the current buffer occupancy **650(z)** equal to the minimum of the buffer size **174** and the potential buffer occupancy.

(120) Subsequently, the subsequence iterator **250(i)** sets the minimum buffer occupancy **652(z)** equal to the minimum of the current buffer occupancy **650(z)** and the minimum buffer occupancy **652(x)** associated with the subsequence encode list **610(x)**. In a complementary fashion, the subsequence iterator **250(i)** sets the maximum buffer occupancy **654(z)** equal to the maximum of the current buffer occupancy **650(z)** and the maximum buffer occupancy **654(x)** associated with the subsequence encode list **610(x)**.

(121) In alternative embodiments, the subsequence iterator **250** may compute the current buffer occupancy **650**, the minimum buffer occupancy **652**, and the maximum buffer occupancy **654** in any technically feasible fashion and at any level of granularity. For instance, in some embodiments, the subsequence iterator **250(i)** computes the current buffer occupancy **650** at a frame-level of granularity based on the following continuous-time function (1):

$$(122) \quad O(t) = O_0 + tR - \sum_{i=0}^i \text{Math. } S_i \quad (1)$$

(123) In function (1), $O(t)$ is a buffer occupancy at the time t , $O_{\text{sub}.0}$ is an initial buffer occupancy at a time of 0, t is a time; R is an average frame rate, ΔT is the time interval between consecutive displayed time (i.e., the inverse of the frame rate), and $S_{\text{sub}.i}$ is the compressed frame size of the $i_{\text{sup}.th}$ frame.

(124) As shown, the filtering engine **730** generates a compliant media encode set **740** based on the media encode set **720** and the transmission buffer constraint **170**. More precisely, the filtering engine **730** initializes the compliant media encode set **740** to an empty set. Subsequently, for each media encode point **270(z)** included in the media encode set **720**, the filtering engine **730** determines whether the associated encoded media sequence **630(z)** is compliant with the transmission buffer constraint **170**. If the filtering engine **730** determines that the encoded media sequence **630(z)** is compliant with the transmission buffer constraint **170**, then the filtering engine **730** adds the media encode point **270(z)** to the compliant media encode set **740**. By contrast, if the filtering engine **730** determines that the encoded media sequence **630(z)** is not compliant with the transmission buffer constraint **170**, then the filtering engine **730** does not add the media encode point **270(z)** to the compliant media encode set **740**.

(125) The filtering engine **730** may determine whether the encoded media sequence **630(z)** is compliant with the transmission buffer constraint **170** in any technically feasible fashion. For instance, in some embodiments, the filtering engine **730** computes a buffer difference between the maximum buffer occupancy **654(z)** and the minimum buffer occupancy **652(z)**. If the buffer difference does not exceed the buffer size **174**, then the filtering engine **730** determines that the encoded media sequence **630(z)** is compliant with the transmission buffer constraint **170**. Otherwise, the filtering engine determines that the encoded media sequence **630(z)** is not compliant with the transmission buffer constraint **170**.

(126) The global convex hull generator **750** performs convex hull operations across the media encode points **270** included in the compliant media encode set **740** to generate the global convex hull **260(i)**. In general, each of the media encode points **270** included in the global convex hull **260(i)** minimizes the average bitrate **632** for a different overall distortion level **636**. The global convex hull generator **750** may perform any number and type of convex hull operations in any technically feasible fashion to generate the global convex hull **260(i)**. Persons skilled in the art will understand that many techniques for generating convex hulls are well known in the field of mathematics, and all such techniques may be implemented to generate the global convex hull **260(i)**.

(127) For instance, in some embodiments, for each of the media encode points **270** included in the compliant media encode set **740**, the global convex hull generator **750** generates a corresponding plotted media encode point **270** in a graph (not shown). The global convex hull generator **750** plots a given media encode point **270** by locating the average bitrate **632** along the bitrate axis **510** and overall distortion level **636** along the distortion axis **520**. The global convex hull generator **750** then

evaluates the plotted media encode points **270** to determine the global convex hull **260(i)**.

(128) More specifically, the global convex hull generator **750** identifies the plotted media encode points **270** that form a boundary where all the plotted media encode points **270** reside on one side of the boundary and also are such that connecting any two consecutive identified plotted media encode points **270** with a straight line leaves all remaining plotted media encode points **270** on the same side. Notably, if multiple plotted media encode points **270** are within a relatively smaller distance of each other relative to the graph, then the global convex hull generator **750** selects the media encode point **270** having the highest current buffer occupancy **650**. The convex hull **260(i)** includes the set of the identified media encode points **270**.

(129) FIG. **8** illustrates an exemplary global convex hull **260(S)** and an exemplary optimized encoded media sequence **180** that are generated by the buffer-based encoding application **140** of FIG. **2**, according to various embodiments of the present invention. More specifically, the subsequence iterator **250(S)** generates the global convex hull **260(S)** of media encode points **270**, where each media encode point is compliant with the transmission buffer constraint **170** and minimizes the average bitrate **632** for a different overall distortion level **636**. For explanatory purposes, FIG. **8** depicts the media encode points **270** plotted as part of a graph **900** that includes, without limitation, the bitrate axis **510** and the distortion axis **520**.

(130) As described previously in conjunction with FIG. **2**, the selection engine **290** determines the optimized encoded media sequence **180** based on the global convex hull **260(S)** that is associated with all the subsequences **132(0)**-**132(S)**. More specifically, the selection engine **290** compares the average bitrates **632** of the media encode points **270** included in the global convex hull **260(S)** and determines that the optimized media encode point **270(W)** has the highest average bitrate **632**. The selection engine **290** then sets the optimized encoded media sequence **180** equal to the encoded media sequence **630(W)** included in the media encode point **270(W)**.

(131) As persons skilled in the art will recognize, the optimized encoded media sequence **180** has a higher overall quality level than any other encoded media sequence **630** associated with the global convex hull **260(S)**. Consequently, the optimized media sequence **180** is an encoded version of the media title that both complies with the transmission buffer constraint **170** and optimizes the overall visual quality during playback of the media title on an endpoint device. In alternative embodiments, the selection engine **290** may select the optimized encoded media sequence **180** based on the global convex hull **260(S)** and any number and type of criteria associated with any number and type of media metrics in any technically feasible fashion.

(132) As shown, the optimized encoded media sequence **180** includes, without limitation, a 960×540 version of the subsequence **132(0)** encoded at QP=40, followed by a 640×360 version of the subsequence **132(1)** encoded at QP=11, followed by a 1280×720 version of the subsequence **132(2)** encoded at QP=47, followed by a 960×540 version of the subsequence **132(3)** encoded at QP=26, and ending with a 640×360 version of the subsequence **132(4)** encoded at QP=2.

Advantageously, as illustrated in FIG. **8**, each of the resolution **314** and the rate control value **316** may vary across the subsequences **132** included in the optimized encoded media sequence **180**.

(133) For explanatory purposes only, the constraint type **176** associated with the exemplary global convex hull **260(S)** is CRB VBV. Consequently, after the selection engine **290** determines the optimized encoded media sequence **180**, the selection engine **290** computes the optimized pre-buffering time **182**. As shown, the selection engine **290** sets an initial buffer occupancy equal to the difference between the buffer size **174** and the minimum buffer occupancy **652(W)** included in the media encode point **270(W)**. The selection engine **290** then sets the optimized pre-buffering time **182** equal to the quotient of the initial buffer occupancy and the transmission rate **172**. In alternative embodiments, the selection engine **190** may compute the optimized pre-buffering time **182** and/or an initial buffer occupancy in any technically feasible fashion.

(134) FIGS. **9A-9B** set forth a flow diagram of method steps for encoding a source media title, according to various embodiments of the present invention. Although the method steps are

described with reference to the systems of FIGS. 1-8, persons skilled in the art will understand that any system configured to implement the method steps, in any order, falls within the scope of the present invention.

(135) As shown, a method **900** begins a step **902**, where the subsequence analyzer **130** partitions the source media sequence **122** into the subsequences **132**. At step **904**, the buffer-based encoding application **140** encodes each of the subsequences **132** across each of the encoding points **212** included in the encoding point list **210** to generate the encoded subsequences **430**. In alternative embodiments, each subsequence **132(x)** may be associated with a different encoding point list **210(x)**. To generate the encoded subsequence **430** associated with both a particular subsequence **132(x)** and a particular encoding point **212(y)**, the buffer-based encoding application **140** encodes the selected subsequence **132** at the associated resolution **314** and the associated rate control value **316**.

(136) At step **906**, for each encoded subsequence **430**, the buffer-based encoding application **140** computes the bitrate **432**, the quality score **434**, and the distortion level **436**. As part of step **906**, the buffer-based encoding application **140** generates a new subsequence encode point **222** that includes, without limitation, the encoding point **212**, the encoded subsequence **430**, the bitrate **432**, the quality score **434**, and the distortion level **436**.

(137) At step **908**, for each subsequence **132(x)**, the convex hull generator **230** generates a different convex hull **240(x)** based on the associated subsequence encode points **222**. Notably, the resolutions **314** and the rate control values **316** may vary between the subsequence encode points **222** included in the convex hull **240(x)**. At step **910**, the buffer-based encoding application **140** sets a previous global convex hull **260** to include a single initialization media encode point **270(-1)** and selects the convex hull **240(0)**. The initialization media encode point **270(-1)** includes, without limitation, the empty subsequence encode list **610**, the NULL encoded media sequence **630**, the current buffer occupancy **650** of the buffer size **174**, the minimum buffer occupancy **652** of the buffer size **174**, and the maximum buffer occupancy **654** of the buffer size **174**.

(138) At step **912**, the combination engine **710** generates the media encode set **720** based on the previous global convex hull **260** and the selected convex hull **240**. First, the combination engine **710** initializes the current media encode set **720** to an empty set. Subsequently, for each media encode point **270(x)** included in the previous global convex hull **260**, the combination engine **710** adds P new media encode points **270** to the media encode set **720**, where P is the total number of subsequence encode points **222** included in the selected convex hull **240**.

(139) To generate the new media encode point **270** associated with both the subsequence encode point **222(y)** and the media encode point **270(x)**, the combination engine **710** appends the subsequence encode point **222(y)** to the subsequence encode list **610(x)** included in the media encode point **270(x)** to generate a new subsequence encode list **610**. The new subsequence encode list **610** is included in the new media encode point **1370**. The combination engine **710** then generates a new encoded media sequence **630** included in the media encode point **270** based on the new subsequence encode list **610**. Finally, the combination engine **710** computes the average bitrate **632**, the overall distortion level **636**, the current buffer occupancy **650**, the minimum buffer occupancy **652**, and the maximum buffer occupancy **654** included in the new media encode point **270** based on the new encoded media sequence **630** and the transmission buffer constraint **170**.

(140) At step **914**, the filtering engine **730** performs filtering operations on the media encode set **720** based on the transmission buffer constraint **170** to generate the compliant media encode set **740**. More precisely, the filtering engine **730** removes the media encode points **270** included in the media encode set **740** for which the difference between the maximum buffer occupancy **654** and the minimum buffer occupancy **652** is greater than the buffer size **174**. At step **916**, the global convex hull generator **750** generates the current global convex hull **260** based on the compliant media encode set **740**.

(141) At step **918**, the subsequence iterator **250** determines whether the selected convex hull **240** is

the final convex hull **240(S)** and, consequently, is associated with the final subsequence **132(S)** included in the source media sequence **122**. If, at step **918**, the subsequence iterator **250** determines that the selected convex hull **240** is not the final convex hull **240(S)**, then the method **900** proceeds to step **920**. At step **920**, the subsequence iterator **250** selects the next convex hull **240** and sets the previous global convex hull **260** equal to the current global convex hull **260**. The method **900** then returns to step **912**, where the subsequence iterator **250** generates a new current global convex hull **260** based on the previous global convex hull **260** and the selected convex hull **240**.

(142) If, however, at step **918**, the subsequence iterator **250** determines that the selected convex hull **240** is the final convex hull **240(S)**, then the method **900** proceeds directly to step **922**. At step **922**, the selection engine **290** selects the media encode point **270** having the highest average bitrate **632** of the media encode points **270** included in the current global convex hull **260**. At step **924**, the selection engine **290** sets the optimized encoded media sequence **180** equal to the encoded media sequence **630** included in the selected media encode point **270**.

(143) At step **926**, the selection engine **290** determines whether the transmission buffer constraint **170** is associated with a configurable pre-buffering time. If, at step **926**, the selection engine **290** determines that the transmission buffer constraint **170** is not associated with a configurable pre-buffering time, then the method **900** proceeds to step **928**. At step **928**, the selection engine **290** transmits the optimized encoded media sequence **180** to the content delivery network **190** for streaming to endpoint devices. The method **900** then terminates.

(144) If, however, at step **926**, the selection engine **290** determines that the transmission buffer constraint **170** is associated with a configurable pre-buffering time, then the method **900** proceeds directly to step **930**. At step **930**, the selection engine **290** computes the optimized pre-buffering time **182** based on the selected media encode point **270** and the transmission buffer constraint **170**. At step **932**, the selection engine **290** transmits the optimized encoded media sequence **180** and the optimized pre-buffering time **182** to the content delivery network **190** for streaming to endpoint devices. The method **900** then terminates.

(145) In sum, the disclosed techniques enable efficient and optimal encoding of source media sequences to generate optimized encoded media sequences that are compliant with transmission buffer constraints. A subsequence-based encoding subsystem includes, without limitation, a subsequence analyzer and a buffer-based encoding application. Initially, the subsequence analyzer partitions a source media sequence into multiple subsequences. For each subsequence, the buffer-based encoding subsystem generates multiple encoded subsequences based on an encoding list. The encoding list includes, without limitation, any number of encoding points, where each encoding point specifies a different combination of a resolution and a rate control value.

(146) For each encoded subsequence, the buffer-based encoding application generates a different subsequence encode point. Each subsequence encode point includes, without limitation, the associated encoded subsequence, the associated encoding point, a bitrate of the encoded subsequence, a quality score of the encoded subsequence, and a distortion level of the encoded subsequence. Subsequently, for each subsequence, a convex hull generator generates a convex hull that includes a subset of the subsequence encode points associated with the subsequence. In general, for a given subsequence, the subsequence encode points included in the associated convex hull minimize the bitrate for different distortion levels.

(147) A subsequence iterator then iteratively processes the subsequences in a temporal playback order to generate a different global convex hull of media encode points for each subsequence. For the n^{th} subsequence, the subsequence iterator computes an n^{th} global convex hull of media encode points based on a transmission buffer constraint, an $(n-1)^{\text{th}}$ global convex hull of media encode points, and the n^{th} convex hull of subsequence encode points associated with the n^{th} subsequence. The transmission buffer constraint specifies, without limitation, a transmission rate, a buffer size, and a constraint type. Each media encode point includes, without limitation, a subsequence encode list, an encoded media sequence, an average bitrate, an overall

distortion level, a current buffer occupancy, a minimum buffer occupancy, and a maximum buffer occupancy.

(148) Each media encode point in the $n^{\text{sup.th}}$ global convex hull is associated with a portion of the media title that spans from the $0^{\text{sup.th}}$ subsequence to the $n^{\text{sup.th}}$ subsequence. To generate the $n^{\text{sup.th}}$ global convex hull, the subsequence iterator combines each of the V media encode points included in the $(n-1)^{\text{sup.th}}$ global convex hull with each of the A subsequence encode points included in the $n^{\text{sup.th}}$ convex hull to generate a media encode set that includes $(V \times A)$ media encode points. The subsequence iterator then performs filtering operations on the media encode set based on the transmission buffer constraint, the minimum buffer occupancies, and the maximum buffer occupancies to generate a compliant media encode set. Finally, the subsequence iterator generates the $n^{\text{sup.th}}$ global convex hull that is a subset of the compliant media encode set based on the average bitrates and the overall distortion levels. In general, for the portion of the media title spanning from the $0^{\text{sup.th}}$ subsequence to the $n^{\text{sup.th}}$ subsequence, the media encode points included in the $n^{\text{sup.th}}$ global convex hull minimize the average bitrate for different overall distortion levels.

(149) After generating the final global convex hull that spans the entire media title, a selection engine included in the buffer-based encoding application selects one of media encode points included in the final global convex hull as the optimized media encode. The optimized media encode point has an average bitrate that is higher than the average bitrates of all other media encode points included in the final global convex hull. The selection engine then sets the optimized encoded media sequence equal to the encoded media sequence included in the optimized media encode point. If the constraint type is associated with a configurable pre-buffering time, then the selection engine computes an optimized pre-buffering time based on the buffer size and the minimum buffer occupancy included in the optimized media encode point. Finally, the selection engine transmits the optimized encoded media sequence and, optionally, the optimized pre-buffering time to a content delivery network for distribution to endpoint devices.

(150) At least one advantage and technological advancement of the disclosed techniques is that encoded media sequences are incrementally generated in the playback order based on sets of individually encoded subsequences and the transmission buffer constraint. Aggregating subsequences encoded using different encoding points to generate the encoded media sequences reduces the encoding inefficiencies typically associated with prior-art monolithic encoding techniques. Further, excluding encoded media sequences that are not compliant with the transmission buffer constraint from the global convex hulls efficiently ensures that the optimized encoded media sequence is compliant with the transmission buffer constraint. As a result, the overall visual quality of the optimized encoded media sequence is increased without increasing the likelihood of playback interruptions attributable to transmission buffer limitations during playback of the media title on endpoint devices. These technical advantages provide one or more technological advancements over the prior art. 1. In some embodiments, a computer implemented method comprises generating a first convex hull of subsequence encode points based on a first set of encoding points and a first subsequence associated with a media title; generating a first global convex hull of media encode points based on a transmission buffer constraint, the first convex hull, and a second global convex hull of media encode points associated with a portion of the media title that occurs before the first subsequence in a playback order for the media title; selecting a first media encode point included in the first global convex hull based on a media metric; and determining a first encoded media sequence based on the first media encode point, wherein the first encoded media sequence is subsequently streamed to an endpoint device during playback of the media title. 2. The computer-implemented method of clause 1, wherein the transmission buffer constraint comprises a constant bitrate video buffer verifier (VBV) constraint or a variable bitrate VBV constraint. 3. The computer-implemented method of clauses 1 or 2, further comprising computing a pre-buffering time based on a minimum buffer occupancy associated with the first

encoded media sequence and a buffer size associated with the transmission buffer constraint. 4. The computer-implemented method of any of clauses 1-3, wherein generating the first global convex hull comprises for each combination of a media encode point included in the second global convex hull and a subsequence encode point included in the first convex hull, generating a different media encode point included in a set of media encode points; performing one or more filtering operations on the set of media encode points based on the transmission buffer constraint to generate a set of compliant media encode points; and performing one or more convex hull operations across the set of compliant media encode points to generate the first global convex hull. 5. The computer-implemented method of any of clauses 1-4, wherein the media metric comprises a bitrate, and selecting the first media encode point comprises determining that a first average bitrate associated with the first media encode point is greater than average bitrates associated with all other media encode points included in the first global convex hull. 6. The computer-implemented method of any of clauses 1-5, wherein generating the first global convex hull comprises generating a second encoded media sequence based on a second media encode point included in the second global convex hull and a first subsequence encode point included in the first convex hull; computing a minimum buffer occupancy and a maximum buffer occupancy based on the second encoded media sequence and a transmission rate associated with the transmission buffer constraint; determining that a difference between the maximum buffer occupancy and the minimum buffer occupancy does not exceed a buffer size associated with the transmission buffer constraint; and in response to determining that the difference does not exceed the buffer size, adding a second media encode point associated with the second encoded media sequence to the first global convex hull. 7. The computer-implemented method of any of clauses 1-6, wherein generating the first convex hull comprises generating a first set of subsequence encode points based on the first set of encoding point and the first subsequence; determining a region that includes the first set of subsequence encode points; identifying a boundary associated with the region, wherein no subsequence encode points included in the first set of subsequence encode points are located on a first side of the boundary; and discarding all subsequence encode points included in the first set of subsequence encode points that are not located along the boundary. 8. The computer-implemented method of any of clauses 1-7, wherein each encoding point included in the first set of encoding points specifies a different combination of a resolution and a rate control value. 9. The computer-implemented method of any of clauses 1-8, further comprising partitioning a source media sequence associated with the media title into a plurality of subsequences based on a plurality of shot changes, wherein the plurality of subsequences includes the first subsequence. 10. The computer-implemented method of any of clauses 1-9, wherein the media title comprises at least one of video content and audio content. 11. In some embodiments, one or more non-transitory computer readable media include instructions that, when executed by one or more processors, cause the one or more processors to perform the steps of encoding a first subsequence associated with a media title across a first set of encoding points to generate a first set of subsequence encode points, wherein each subsequence encode point is associated with a different encoded subsequence; performing one or more convex hull operations across the first set of subsequence encode points to compute a first convex hull of subsequence encode points based on at least two media metrics; generating a first global convex hull of media encode points based on a transmission buffer constraint, the first convex hull, and a second global convex hull of media encode points associated with one or more subsequences associated with the media title that occur before the first subsequence in a playback order for the media title; selecting a first media encode point included in the first global convex hull based on a first media metric; and determining a first encoded media sequence based on the first media encode point, wherein the first encoded media sequence is subsequently streamed to an endpoint device during playback of the media title. 12. The one or more non-transitory computer readable media of clause 11, wherein the transmission buffer constraint specifies at least one of a transmission rate and a buffer size. 13. The one or more non-

transitory computer readable media of clauses 11 or 12, further comprising computing a pre-buffering time based on a minimum buffer occupancy associated with the first encoded media sequence and a buffer size associated with the transmission buffer constraint. 14. The one or more non-transitory computer readable media of any of clauses 11-13, wherein generating the first global convex hull comprises generating a second media encode point based on a third media encode point included in the second global convex hull and a first subsequence encode point included in the first convex hull; generating a fourth media encode point based on the third media encode point and a second subsequence encode point included in the first convex hull; adding the second media encode point but not the fourth media encode point to a set of compliant encode points based on the transmission buffer constraint; and performing one or more convex hull operations across the set of compliant media encode points to generate the first global convex hull. 15. The one or more non-transitory computer readable media of any of clauses 11-14, wherein the first media metric comprises a bitrate, and selecting the first media encode point comprises determining that a first average bitrate associated with the first media encode point is greater than average bitrates associated with all other media encode points included in the first global convex hull. 16. The one or more non-transitory computer readable media of any of clauses 11-15, wherein generating the first global convex hull comprises generating a second encoded media sequence based on a second media encode point included in the second global convex hull and a first subsequence encode point included in the first convex hull; computing a minimum buffer occupancy and a maximum buffer occupancy based on the second encoded media sequence and a transmission rate associated with the transmission buffer constraint; determining that a difference between the maximum buffer occupancy and the minimum buffer occupancy does not exceed a buffer size associated with the transmission buffer constraint; and in response to determining that the difference does not exceed the buffer size, adding a second media encode point associated with the second encoded media sequence to the first global convex hull. 17. The one or more non-transitory computer readable media of any of clauses 11-16, wherein generating the first global convex hull comprises generating a set of compliant media encode points based on the transmission buffer constraint, the first convex hull, and the second global convex hull; determining a region that includes the set of compliant media encode points; identifying a boundary associated with the region, wherein no compliant media encode points included in the set of compliant media encode points are located on a first side of the boundary; and discarding all compliant media encode points included in the set of compliant media encode points that are not located along the boundary. 18. The one or more non-transitory computer readable media of any of clauses 11-17, wherein the first encoded media sequence includes a first encoded subsequence and a second encoded subsequence, and wherein the first encoded subsequence is associated with a first encoding point included in the first set of encoding points and the second encoded subsequence is associated with a second encoding point included in the first set of encoding points that is not equal to the first encoding point. 19. The one or more non-transitory computer readable media of any of clauses 11-18, wherein the first subsequence associated with the media title comprises a shot sequence. 20. In some embodiments, a system, comprises one or more memories storing instructions; and one or more processors that are coupled to the one or more memories and, when executing the instructions, are configured to partition a source media sequence associated with a media title into a plurality of subsequences based on a plurality of shot changes; generate a first convex hull of subsequence encode points based on a first set of encoding points and a first subsequence included in the plurality of subsequences, wherein each subsequence encode point is associated with a different encoded subsequence; generate a first global convex hull of media encode points based on a transmission buffer constraint, the first convex hull, and a second global convex hull of media encode points associated with a portion of the media title that occurs before the first subsequence in a playback order for the media title; select a first media encode point included in the first global convex hull based on a media metric; and determine a first encoded media sequence based on the first media

encode point, wherein the first encoded media sequence is subsequently streamed to an endpoint device during playback of the media title.

(151) Any and all combinations of any of the claim elements recited in any of the claims and/or any elements described in this application, in any fashion, fall within the contemplated scope of the present invention and protection.

(152) The descriptions of the various embodiments have been presented for purposes of illustration, but are not intended to be exhaustive or limited to the embodiments disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the described embodiments.

(153) Aspects of the present embodiments may be embodied as a system, method or computer program product. Accordingly, aspects of the present disclosure may take the form of an entirely hardware embodiment, an entirely software embodiment (including firmware, resident software, micro-code, etc.) or an embodiment combining software and hardware aspects that may all generally be referred to herein as a “module” or “system.” In addition, any hardware and/or software technique, process, function, component, engine, module, or system described in the present disclosure may be implemented as a circuit or set of circuits. Furthermore, aspects of the present disclosure may take the form of a computer program product embodied in one or more computer readable medium(s) having computer readable program code embodied thereon.

(154) Any combination of one or more computer readable medium(s) may be utilized. The computer readable medium may be a computer readable signal medium or a computer readable storage medium. A computer readable storage medium may be, for example, but not limited to, an electronic, magnetic, optical, electromagnetic, infrared, or semiconductor system, apparatus, or device, or any suitable combination of the foregoing. More specific examples (a non-exhaustive list) of the computer readable storage medium would include the following: an electrical connection having one or more wires, a portable computer diskette, a hard disk, a random access memory (RAM), a read-only memory (ROM), an erasable programmable read-only memory (EPROM or Flash memory), an optical fiber, a portable compact disc read-only memory (CD-ROM), an optical storage device, a magnetic storage device, or any suitable combination of the foregoing. In the context of this document, a computer readable storage medium may be any tangible medium that can contain or store a program for use by or in connection with an instruction execution system, apparatus, or device.

(155) Aspects of the present disclosure are described above with reference to flowchart illustrations and/or block diagrams of methods, apparatus (systems) and computer program products according to embodiments of the disclosure. It will be understood that each block of the flowchart illustrations and/or block diagrams, and combinations of blocks in the flowchart illustrations and/or block diagrams, can be implemented by computer program instructions. These computer program instructions may be provided to a processor of a general purpose computer, special purpose computer, or other programmable data processing apparatus to produce a machine. The instructions, when executed via the processor of the computer or other programmable data processing apparatus, enable the implementation of the functions/acts specified in the flowchart and/or block diagram block or blocks. Such processors may be, without limitation, general purpose processors, special-purpose processors, application-specific processors, or field-programmable gate arrays.

(156) The flowchart and block diagrams in the figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods and computer program products according to various embodiments of the present disclosure. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of code, which comprises one or more executable instructions for implementing the specified logical function(s). It should also be noted that, in some alternative implementations, the functions noted in the block may occur out of the order noted in the figures. For example, two blocks shown in succession may, in fact, be

executed substantially concurrently, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts, or combinations of special purpose hardware and computer instructions.

(157) While the preceding is directed to embodiments of the present disclosure, other and further embodiments of the disclosure may be devised without departing from the basic scope thereof, and the scope thereof is determined by the claims that follow.

Claims

1. A method implemented by an endpoint device, the method comprising: receiving a first encoded media sequence of a media title that has been encoded based on an encode point included in a convex hull of encode points, wherein the convex hull of encode points is generated based on at least two other convex hulls of encode points associated with different portions of the media title and is filtered based on one or more transmission buffer constraints associated with the endpoint device; storing the first encoded media sequence in a transmission buffer included in the endpoint device; and decoding for playback at least a portion of the first encoded media sequence stored in the transmission buffer.
2. The method of claim 1, wherein the encode point minimizes an average bitrate subject to the one or more transmission buffer constraints.
3. The method of claim 1, wherein the one or more transmission buffer constraints include at least one transmission buffer constraint limiting a variation in bits used to encode frames included in the first encoded media sequence.
4. The method of claim 1, wherein the one or more transmission buffer constraints include at least one transmission buffer constraint associated with a pre-buffering time.
5. The method of claim 4, further comprising receiving the pre-buffering time.
6. The method of claim 4, wherein the pre-buffering time is based on an initial buffer occupancy and a transmission rate.
7. The method of claim 1, wherein the one or more transmission buffer constraints include at least one of a constant bitrate video buffer verifier (VBBV) constraint or a variable bitrate VBBV constraint.
8. The method of claim 1, wherein the first encoded media sequence comprises a plurality of subsequences, and at least two of the subsequences are encoded based on different encoding parameters and subject to the one or more transmission buffer constraints.
9. The method of claim 1, wherein the first encoded media sequence is associated with a media title comprising at least one of video content or audio content.
10. The method of claim 1, further comprising: receiving a second encoded media sequence that has been encoded based on at least a first encoding parameter and to comply with the one or more transmission buffer constraints, wherein the at least a first encoding parameters is different than at least one encoding parameter used to encode the first encoded media sequence; and storing the second encoded media sequence in the transmission buffer; and decoding for playback at least a portion of the second encoded media sequence.
11. The method of claim 1, wherein the at least two other convex hulls include a first convex hull of encode points corresponding to a first subsequence of the media title, and a second convex hull of encode points corresponding to a portion of the media title that occurs before the first subsequence in a playback order for the media title.
12. One or more non-transitory computer readable media including instructions that, when executed by one or more processors included in an endpoint device, cause the one or more

processors to perform the steps of: receiving a first encoded media sequence of a media title that has been encoded based on an encode point included in a convex hull of encode points, wherein the convex hull of encode points is generated based on at least two other convex hulls of encode points associated with different portions of the media title and is filtered based on one or more transmission buffer constraints associated with the endpoint device; storing the first encoded media sequence in a transmission buffer included in the endpoint device; and decoding for playback at least a portion of the first encoded media sequence stored in the transmission buffer.

13. The one or more non-transitory computer readable media of claim 12, wherein the encode point minimizes an average bitrate subject to the one or more transmission buffer constraints.

14. The one or more non-transitory computer readable media of claim 12, wherein the one or more transmission buffer constraints include at least one of a transmission buffer constraint associated with a pre-buffering time or a transmission buffer constraint limiting a variation in bits used to encode frames included in the first encoded media sequence.

15. The one or more non-transitory computer readable media of claim 14, further comprising receiving the pre-buffering time.

16. The one or more non-transitory computer readable media of claim 14, wherein the pre-buffering time is based on an initial buffer occupancy and a transmission rate.

17. The one or more non-transitory computer readable media of claim 12, wherein the first encoded media sequence comprises a plurality of subsequences, and at least two of the subsequences are encoded based on different encoding parameters and subject to the one or more transmission buffer constraints.

18. The one or more non-transitory computer readable media of claim 12, the steps further comprising: receiving a second encoded media sequence that has been encoded based on at least a first encoding parameter and to comply with the one or more transmission buffer constraints, wherein the at least a first encoding parameters is different than at least one encoding parameter used to encode the first encoded media sequence; and storing the second encoded media sequence in the transmission buffer; and decoding for playback at least a portion of the second encoded media sequence.

19. The one or more non-transitory computer readable media of claim 12, wherein the first encoded media sequence is downloaded from a content distribution network.

20. The one or more non-transitory computer readable media of claim 12, wherein the first encoded media sequence comprises a real-time video stream.

21. An endpoint device, comprising: one or more memories storing instructions; and one or more processors that are coupled to the one or more memories and, when executing the instructions, are configured to: receive a first encoded media sequence of a media title that has been encoded based on an encode point included in a convex hull of encode points, wherein the convex hull of encode points is generated based on at least two other convex hulls of encode points associated with different portions of the media title and is filtered based on one or more transmission buffer constraints associated with the endpoint device; store the first encoded media sequence in a transmission buffer included in the endpoint device; and decode at least a portion of first encoded media sequence stored in the transmission buffer for playback.
