



US 20250258613A1

(19) **United States**

(12) **Patent Application Publication**
Mylly

(10) **Pub. No.: US 2025/0258613 A1**

(43) **Pub. Date: Aug. 14, 2025**

(54) **CHANNEL OPTIMIZED STORAGE
MODULES**

(60) Provisional application No. 61/915,315, filed on Dec.
12, 2013.

(71) Applicant: **Memory Technologies LLC**, Kirkland,
WA (US)

Publication Classification

(72) Inventor: **Kimmo Juhani Mylly**, Ylojarvi (FI)

(51) **Int. Cl.**
G06F 3/06 (2006.01)

(73) Assignee: **Memory Technologies LLC**, Kirkland,
WA (US)

(52) **U.S. Cl.**
CPC **G06F 3/0619** (2013.01); **G06F 3/061**
(2013.01); **G06F 3/0644** (2013.01); **G06F**
3/0659 (2013.01); **G06F 3/0679** (2013.01);
G06F 3/0688 (2013.01)

(21) Appl. No.: **19/020,614**

(22) Filed: **Jan. 14, 2025**


(57) **ABSTRACT**

Related U.S. Application Data

(63) Continuation of application No. 18/500,095, filed on
Nov. 1, 2023, which is a continuation of application
No. 17/327,153, filed on May 21, 2021, now Pat. No.
11,809,718, which is a continuation of application
No. 16/720,698, filed on Dec. 19, 2019, now Pat. No.
11,023,142, which is a continuation of application
No. 16/545,899, filed on Aug. 20, 2019, now Pat. No.
11,016,678, which is a continuation of application
No. 15/705,559, filed on Sep. 15, 2017, now Pat. No.
10,402,106, which is a continuation of application
No. 14/569,601, filed on Dec. 12, 2014, now Pat. No.
9,766,823.

A storage module includes a set of memories. Each of the
memories in the set of memories may be divided into a set
of portions. A controller is configured to transfer data
between the set of memories and a host connected through
an interface. A set of channels connects the set of memories
to the controller. The controller is also configured to select:
a memory from the set of memories, a portion from the set
of portions for the selected memory, and/or a channel from
the set of channels, e.g., connected to the selected memory,
based upon an identification (ID) associated with the data.
The ID may be separate from the data and a write address of
the data, and the selected memory, the selected portion, and
the selected channel may be used to store the data.

100



100




Fig. 1




Fig. 2




Fig. 3




Fig. 4

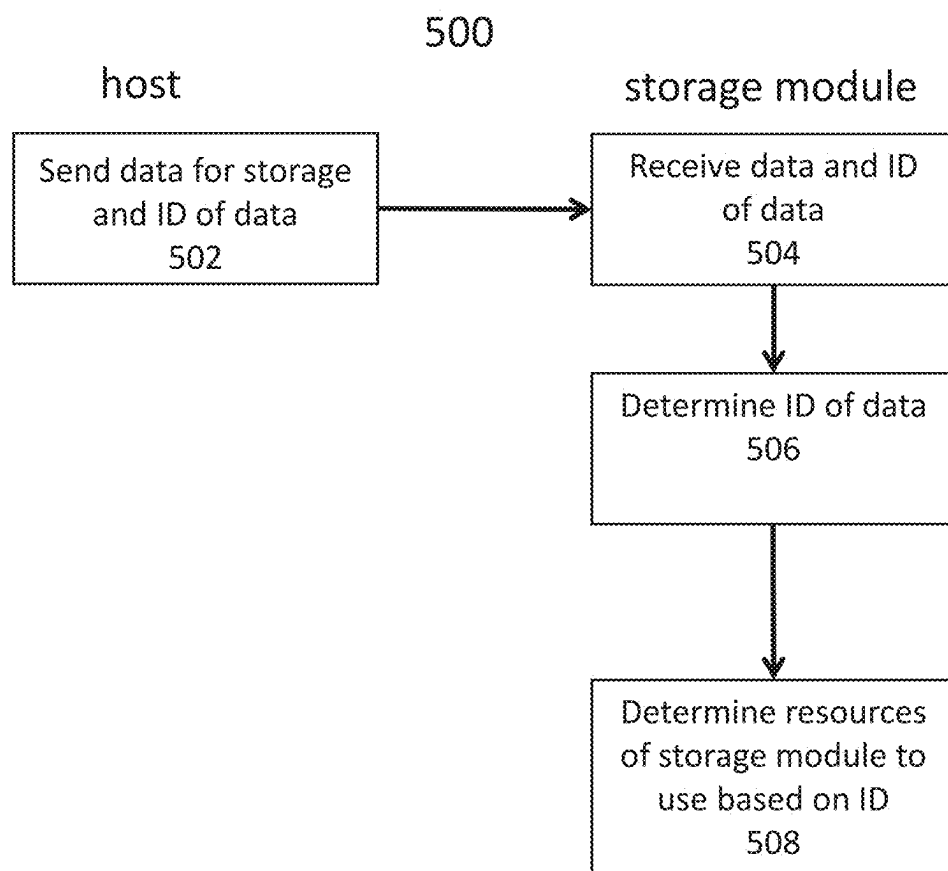


Fig. 5

CHANNEL OPTIMIZED STORAGE MODULES

RELATED APPLICATIONS

[0001] This application is a continuation of co-pending U.S. patent application Ser. No. 18/500,095 filed Nov. 1, 2023, which is a continuation of U.S. patent application Ser. No. 17/327,153 filed May 21, 2021, which is a continuation of U.S. patent application Ser. No. 16/720,698 filed Dec. 19, 2019, which is a continuation of U.S. patent application Ser. No. 16/545,899 filed Aug. 20, 2019, now U.S. Pat. No. 11,016,678, which is a continuation of Ser. No. 15/705,559 filed Sep. 15, 2017, now U.S. Pat. No. 10,402,106, which is a continuation of U.S. patent application Ser. No. 14/569,601 filed Dec. 12, 2014, now U.S. Pat. No. 9,766,823, which claims priority to U.S. Provisional Application No. 61/915,315 filed Dec. 12, 2013 and entitled “Channel Optimized Storage Modules,” all of which are incorporated herein in their entirety. This application is related to U.S. application Ser. No. 12/455,763, filed Jun. 4, 2009, now U.S. Pat. No. 8,874,824, U.S. application Ser. No. 11/176,669, filed Jul. 8, 2005, now U.S. Pat. No. 7,827,370, and U.S. application Ser. No. 13/179,689, filed Jul. 11, 2011, now U.S. Pat. No. 8,560,778, each of which are assigned to the assignee of the present application. Each of these related applications are incorporated herein by reference in their entireties.

BACKGROUND

[0002] Managed storage modules, such as managed NAND storage modules, provide many benefits over using raw memories such as flash NAND memories. Managed storage modules, which typically include a storage controller combined with NAND memory in the case of managed NAND or other types of memory in other cases, provide several benefits to device manufacturers. The storage controller hides the details of the memory (e.g., NAND) and provides the intended interface and other features, such as error-correcting code (ECC) support without the device manufacturers having to implement those features on the host side (e.g., a smartphone or a tablet, an SoC). Additionally, managed storage modules allow new advanced features to be implemented in the storage controller without the host necessarily having to be aware that the features exist. The advanced features may either be activated or not by the storage controller depending on whether the host supports the features. Thus, managed storage modules improve backwards compatibility.

[0003] Examples of managed storage modules (e.g., managed storage devices), and in particular managed NAND storage modules, include embedded multimedia cards (eMMC), Universal Flash Storage (UFS), solid-state drive (SSD) modules. These modules are used in wide variety of applications like mobile phones, Global positioning system (GPS) devices, media players, PCs, and servers for storing the operating system code, applications, and user data, such as, photos and videos. Along with the data visible to the host device, operational code/firmware (FW) of the storage module itself is stored in the memory of the storage module. Additionally, other important data, which is used to operate the memory module, such as register data and address translation data, may be stored in the memory.

[0004] The increase of interface bandwidth and introduction of caching has improved the sequential read/write and

random write performance of a managed storage module. Queuing commands may also improve the random read performance. However, this depends on a capability of the system to efficiently queue commands and, to some extent, a capability to predict how the storage will be accessed in the future. U.S. Pat. No. 8,560,778, filed Jul. 11, 2011, introduces some methods for informing the storage module to pre-fetch data before a new command is received. The content of this patent is incorporated herein in its entirety.

SUMMARY

[0005] Embodiments of the present disclosure include a storage module including a set of memories including at least one memory, wherein each of the memories in the set of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion. The storage module also includes a controller configured to transfer data between the set of memories and a host connected through an interface. The storage module also includes a set of channels connecting the set of memories to the controller, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects the memory controller to at least one of the memories in the set of memories. The controller is configured to select a first memory from the set of memories, a portion from the set of portions for the first memory, or a first channel from the set of channels connected to the first memory based upon an ID associated with the data, wherein the ID is separate from the data and the write address of the data, and wherein the first memory, the first portion, and the first channel are used to store the data.

[0006] Embodiments of the present disclosure include a storage module comprising a set of memories including at least one memory, wherein each of the memories in the set of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion. A controller is configured to transfer the data between the set of memories and a host through an interface. A set of channels connects the set of memories to the controller, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects the controller to at least one of the memories in the set of memories. The controller is further configured to select at least one of (1) a memory area from the set of memories or (2) a channel from the set of channels, based at least in part upon an ID associated with the data. The at least one of the selected memory area or the selected channel are used to store the data.

[0007] In some embodiments, the selected memory area comprises a selected memory from the set of memories. In some embodiments, the controller is further configured to select a portion from the set of portions of the selected memory.

[0008] In some embodiments, the selected memory area comprises a selected portion from the set of portions of a memory of the set of memories.

[0009] In some embodiments, the set of memories is a set of non-volatile memories.

[0010] In some embodiments, the set of memories is one of NAND, eMMC, UFS, SD card or SSD memory.

[0011] In some embodiments, the set of memories includes a first memory and a second memory.

[0012] In some embodiments, the set of portions for each memory in the set of memories includes a first portion and a second portion.

[0013] In some embodiments, the set of channels includes a first channel and a second channel.

[0014] In some embodiments, each of the memories in the set of memories is a memory chip.

[0015] In some embodiments, each of the portions in the set of portions is a plane of a memory.

[0016] In some embodiments, the controller is further configured to write data with a same ID using a same memory, portion, and channel.

[0017] In some embodiments, the controller is further configured to extract the ID prior to receiving the data.

[0018] In some embodiments, the controller is further configured to extract the ID as part of decoding a write or read command for the data.

[0019] In some embodiments, a storage device may further comprise at least one register containing a number of parallel accesses supported by the storage device, wherein a number of distinct IDs supported by the controller is a same number as the number of parallel accesses supported by the storage device.

[0020] In some embodiments, a storage device may further comprise at least one register containing at least one of an indication of a width of channels in the set of channels, an indication of a number of channels in the set of channels, an indication of a number of memory chips in the set of memories, a number of memory chips per channel, or a number of planes implemented in an individual memory chip.

[0021] In some embodiments, a first memory is selected by the controller and a first channel and a first portion are used as a default.

[0022] In some embodiments, a first portion is selected by the controller and a first memory and a first channel are used as a default.

[0023] Embodiments of the present disclosure include a method comprising receiving data to be stored in a storage device. The storage device comprises a set of memories including at least one memory, wherein each of the memories in the set of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion. The storage device comprises a set of channels connecting the set of memories to the controller, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects a controller of the storage device to at least one of the memories in the set of memories. The method further comprises receiving an ID associated with the data to be stored in the storage device, selecting, by the storage device and based at least in part upon the ID associated with the data, at least one of (1) a memory area from the set of memories or (2) a channel from the set of channels, and utilizing, by the storage device, the at least one of the selected memory area or the selected channel to store the data.

[0024] Embodiments of the present disclosure include a storage device comprising means for receiving data to be stored in the storage device and an ID associated with the data to be stored in the storage device. The storage device also comprises means for controlling the storage device. The storage device comprises a set of memories including at least one memory, wherein each of the memories in the set

of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion. The storage device comprises a set of channels connecting the set of memories to the means for controlling the storage device, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects the means for controlling the storage device to at least one of the memories in the set of memories. The means for controlling the storage device selects, based at least in part upon the ID associated with the data, at least one of (1) a memory area from the set of memories or (2) a channel from the set of channels and utilizes the at least one of the selected memory area or the selected channel to store the data.

[0025] Embodiments of the present disclosure include a host device including an interface configured to couple the host device to a storage device. The storage device includes a set of memories including at least one memory, wherein each of the memories in the set of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion. The storage device includes a set of channels connecting the set of memories to the controller, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects a controller of the storage device to at least one of the memories in the set of memories. The host device also includes a host controller configured to detect a number of possible parallel read accesses associated with data to be stored in the storage device, determine, based on the number of possible parallel read accesses associated with the data, an ID associated with the data; and transmit, to the storage device, the ID associated with the data.

[0026] In some embodiments, the host controller is configured to detect the number of possible parallel read accesses associated with the data by at least reading an indication of the number of possible parallel read accesses associated with the data from a register of the storage device.

[0027] In some embodiments, the host controller is configured to detect the number of possible parallel read accesses associated with the data by at least receiving an indication of the number of possible parallel read accesses associated with the data from the storage device.

BRIEF DESCRIPTION OF THE DRAWINGS

[0028] FIG. 1 depicts an example host device.

[0029] FIG. 2 depicts an embodiment of a storage module.

[0030] FIG. 3 depicts an embodiment of the storage module illustrating additional detail associated with a mass storage.

[0031] FIG. 4 depicts an alternative embodiment of the storage module illustrating additional detail associated with a mass storage.

[0032] FIG. 5 depicts a process flow of a first embodiment of the present disclosure.

[0033] The figures depict embodiments of the present disclosure for purposes of illustration only. One skilled in the art will readily recognize from the following discussion that alternative embodiments of the structures and methods illustrated herein may be employed without departing from embodiments described herein.

DETAILED DESCRIPTION

[0034] The following description is presented to enable a person of ordinary skill in the art to make and use the various embodiments. Descriptions of specific devices, techniques, and applications are provided only as examples. Various modifications to the examples described herein will be readily apparent to those of ordinary skill in the art, and the general principles defined herein may be applied to other examples and applications without departing from the spirit and scope of the various embodiments. Thus, the various embodiments are not intended to be limited to the examples described herein and shown, but are to be accorded the scope consistent with the claims.

[0035] A managed storage module's random read performance may be improved by command queuing. For example, if several commands are be queued together, then the managed storage module may determine which of the commands to immediately execute and which commands may need to wait for other commands to finish. In other words, command queuing allows for efficient use of the hardware resources of the managed storage module by ensuring that a maximum number of commands possible is always being executed simultaneously.

[0036] However, in some cases, adding command queuing may not improve performance because all of the commands in the queue may require access to the same hardware resources (e.g., the commands are all accessing the same block of memory). To minimize the chance of this happening, data stored on the storage module may be organized as efficiently as possible. Various different types of access patterns may occur while accessing a storage module (e.g., a relatively short 4 kB random access followed by a relatively long 1 MB sequential access, a relatively long 1 MB sequential access followed by a relatively short 4 kB random access, etc.). Unless a storage module is aware of how data is going to be accessed in the future, there is little that the storage module can do to properly organize the data to make command queuing effective.

[0037] One possible method to improve the organization of the data is to inform the managed storage module about how data may be accessed in the future when first writing the data to the managed storage module. For example, if two streams of data are each marked as different streams and as being likely accessed in parallel in the future, then the managed storage module may store the two streams of data in such a manner as to improve the chance that command queuing is effective. For instance, the managed storage module may store the two data streams in a manner that minimizes the commonality of resources (e.g., memory chips, planes, or channels) that are used to access each data stream. Some of the embodiments of the present disclosure discussed below implement these types of features.

[0038] A method is introduced to maximize the parallelism inside a memory module, particularly related to future coming reads. By informing the memory module about predicted parallel accessed streams the memory module may better arrange the data favorably for parallel accessed streams.

[0039] FIG. 1 depicts an example host 100, e.g., a host device such as a smartphone device or a tablet device, that may utilize embodiments of the present disclosure. Host 100 includes a touch display 102 that is sensitive to a user's touch based on capacitive or resistive detection. Bus 103 connects touch display 102 to processor 104, which may

include a graphics subsystem that handles the display of graphics and text on touch display 102. Host 100 also includes a number of other components connected to processor 104 through shared bus 106, including system memory 108 (e.g., DRAM), sensors 110 (e.g., accelerometers, gyroscope, GPS), input/output (I/O) 112 (e.g., a speaker, a microphone, or a keyboard), and communications interfaces 114 (e.g., USB, WiFi, Bluetooth, or other wired or wireless interfaces). Processor 104 may also include host controller 118 (which may be alternatively connected to but is separate from processor 104) that interfaces with storage module 120 over bus 122. Storage modules according to embodiments, including storage module 120, include one or more storage devices. Alternatively, host controller 118 may interface with storage module 120 over shared bus 106. Both shared bus 106 and bus 122 may include several bus lines for data, commands, clocking signals, power, reset, etc. An example of the bus lines included in bus 122 is described below with respect to FIG. 2. Battery 116 provides power to above described components through a power supply bus and/or lines (not shown).

[0040] While the use of storage module 120 is shown in the context of a touch sensitive smartphone or tablet, the present disclosure is not limited to use in such devices. Embodiments of the present disclosure may be applied to any electronic device that utilizes storage, e.g., wearable computers such as smartwatches or glasses, televisions, cameras, netbooks, gaming consoles, personal computers, servers, set top boxes, and the like. Additionally, the architecture of host 100 is provided for illustrative purposes only and should not be considered limiting.

[0041] FIG. 2 depicts an exemplary architecture for storage module 120 that may implement embodiments of the present disclosure. Storage module 120 may be a memory storage device contained within a package (e.g. a ball grid array (BGA) package) that is designed to be mounted on a printed circuit board. For example, storage module 120 may be an embedded multimedia card (eMMC) or a Universal Flash Storage (UFS) module. Alternatively, storage module 120 may be a memory storage device contained within a removable card that fits within a slot on the host 100 or a semi-removable device such as an SSD module or PC/server cards/modules (e.g., PCIe cards). Additionally, although storage module 120 is shown as being one self-contained storage device, storage module 120 may also be implemented with a collection of interconnected devices.

[0042] As shown in FIG. 2, storage module 120 includes storage controller 200 for communicating data between mass storage 202 and host 100 (see FIG. 1). Storage controller 200 includes control circuit 204 for controlling the operation of storage controller 200. Control circuit 204 may be connected to RAM 214 over bus 213 for storing operating information and/or for temporary storage. Storage controller 200 also includes clock generation circuit 206 for generating an internal clocking signal on internal clock line 207, receiver circuit 208 for receiving data and commands from host controller 118 (see FIG. 1), transmitter circuit 210 for transmitting data and status information to host controller 118 (see FIG. 1), and registers 212 for storing information and settings relating to the operation of storage module 120, including information related to the generation of the internal clocking. Control circuit 204 may use bus 211 to access or write information to registers 212. Storage module 120 communicates with host controller 118 through data out line

215b and data terminal **215a**, which may provide data and status information, and data in line **216b** and data terminal **216a**, which may provide data, commands, and status information.

[0043] Storage module **120** also includes reference clock line **218b** and reference clock terminal **218a** that provide a reference clock signal to clock generating circuit **206**, and power line **220b** and power terminal **220a** that provide power to storage controller **200** and mass storage **202**. While the above lines and terminals are shown to be single lines and terminals in FIG. 2, each line and terminal may be made up of multiple lines and terminals. For example, power terminal **220a** may include multiple terminals associated with multiple lines of power line **220b** that each individually provide power to the different components (e.g., mass storage **202** and storage controller **200**). As another example, data out line **215b** and data out terminal **215a** or data in line **216b** and data in terminal **216a** may be implemented using two lines (e.g., a differential pair or a 2-bit wide bus) connected to two terminals. Bus **222** allows for storage controller **200** to read data from and write data to mass storage **202**.

[0044] Storage module **120** also includes mass storage **202**, which includes one or more memory blocks on one or more memory planes/banks, which may be on one or more chips having memory circuits or cells for storing one or more bits of information. For example, mass storage **202** may be implemented with a non-volatile memory such as NAND flash memory having memory cells/circuits (e.g., NAND cells) each capable of storing one bit (single-level cell) or multiple bits (multi-level cell) of data. Other forms of non-volatile memory may also be used without departing from the present disclosure. Mass storage **202** may be physically and/or logically divided. For example, mass storage **202** may be implemented as a single chip. Alternatively, mass storage **202** may be implemented with several discrete chips that are connected together in a single package (as shown in FIG. 2) or, alternatively, separately packaged and externally connected together. In some examples, mass storage **202** may also be divided up into planes, banks, blocks, and/or pages. Storage controller **200** is connected to mass storage **202** through bus **222**, which allows for storage controller **200** to read data from, and write data to, mass storage **202**.

[0045] RAM **214** is present in some embodiments of the present disclosure; the storage controller **200** may use RAM **214** to store operating information (e.g., operating code and/or state information). For example, RAM **214** may store a translation table that describes how logical addresses are mapped to physical addresses of mass storage **202**. When RAM **214** is not implemented within storage module **120** or is only implemented with limited size within storage module **120**, in some cases, storage controller **200** may request and use a portion of system memory **108** of host **100** (see FIG. 1) in place of RAM **214** or together with any implemented RAM **214**, as described in U.S. patent application Ser. No. 12/455,763, filed Jun. 4, 2009, which is incorporated by reference in its entirety.

[0046] Clock generation circuit **206** may be implemented with a circuit that is capable of generating a clock signal. For example, clock generation circuit **206** may be implemented using common clock recovery and/or generation circuits including PLLs, oscillators, voltage controlled oscillators, delay locked loops, frequency detectors, frequency multi-

pliers/dividers, phase detectors, combinations of these circuits, or any other suitable circuit. Clock generation circuit **206** may also rely on other components, such as resistors, capacitors, inductors, crystals, or MEMS devices. Clock generation circuit **206** may also be programmable so that it may provide a clocking signal output that varies according to the inputs that it receives. For example, clock generation circuit **206** may be configured to produce a clocking signal of a very high quality (e.g., low jitter) when a reference clock signal is present on reference clock line **218b**. Clock generation circuit **206** may also be configured to produce a clocking signal of a lower quality when a reference clock signal is absent. As other examples, the frequency, duty cycle, jitter, output skew, or propagation delay of the outputted clocking signal may be set according to inputs (e.g., control bits) that are provided to clock generation circuit **206** through bus **205**. In alternative architectures, clock generation circuit **206** may have direct access registers **212** without going through control circuit **204**; in still other embodiments clock generation circuit **206** may have a register internal to itself for storing clock configuration information. While clock generation circuit **206** is shown to be part of storage controller **200**, clock generation circuit **206** may also be implemented external to storage controller **200** without departing from the present disclosure.

[0047] Receiver circuit **208** and transmitter circuit **210** receive the internal clock signal on internal clock line **207** so that storage module **120** may transfer data to host **100** at higher rates than without a clock signal. In another embodiment, internal clock line **207** provides the internal clock signal to the receiver circuit **208**, but not to the transmitter circuit **210**. In yet another embodiment, internal clock line **207** provides the internal clock signal to the transmitter circuit **210**, but not to the receiver circuit **208**.

[0048] Registers **212** store one or more bits of information regarding the operation of storage module **120**, including information regarding the operation of clock generation circuit **206** or other features of storage module **120**. Registers **212** may be implemented as part of storage controller **200**, as part of mass storage **202**, as part of RAM **214**, or as part of some other memory circuit in storage module **120**. The memory used for registers **212** may be any type. For example, registers **212** may be implemented in volatile memory (e.g., SRAM, DRAM, or other volatile memory), non-volatile memory (e.g., flash memory, magnetic memory, resistive memory, or other non-volatile memory), read-only memory (ROM), one time programmable memory, or any combination of these or other types of memory.

[0049] Registers **212** may include several individual registers, e.g., registers **212a-212h** of similar or different sizes. For example, register **212a** may be a 1-byte register while registers **212b-212e** are 1-bit registers and register **212f** is a 4-byte register. Registers **212** may be used to store several specific types of information. In one case, some of registers **212** store read-only information that describes how storage module **120** operates (e.g., supported features) or configuration for storage module **120** to properly operate or operate at different levels of performance (e.g., configurations for different transfer rates). In another case, some of registers **212** store writeable information that configures how storage module **120** operates or what specifications storage module **120** uses to operate. In yet another case, some of registers **212** store information about how storage module **120** is currently operating or the current state of storage module

120. Together, registers **212** may also store all of the different types of information described above along with other types of data. Registers **212** may also be used to implement descriptors, flags, and attributes as described in JEDEC Standard No. 220A for Universal Flash Storage (UFS 1.1), published June 2012, which is incorporated by reference herein in its entirety.

[0050] In one case, registers **212** store information that describes a region of mass storage **202** that is write protected (either permanently or temporarily). For example, register **212f** may define an address range, a block range, a partition, or the like that defines the region. Another register, e.g., register **212g**, may define whether the region is permanently, temporarily, or authenticated write protected. In the case of permanent or temporary, the region is protected as described in U.S. Pat. No. 7,827,370, which is incorporated by reference in its entirety. However, in the case of the region being authenticated write protected, the region may be written/programmed to if authentication of the data to be written is successful. Implementation of this feature is discussed below with respect to various embodiments of the present disclosure.

[0051] In another case, registers **212** may store information that describes the hardware resources available in storage module **120**. For example, registers **212** may store an indication of a width of channels in bus **222** or a number of channels in bus **222**. Similarly, registers **212** may also store an indication of a number of memory chips in mass storage **202**, a number of memory chips per channel, and/or a number of planes implemented in an individual memory chip. Examples of a number (1-N) of channels in bus **222**, a number (1-N) of memory chips in mass storage **202**, a number (1-N) of memory chips associated with a channel, and/or a number (1-N) of planes per memory chip, are further discussed below with respect to FIGS. 3 and 4. Registers **212** may also store an indication of how many parallel data streams may be read or written at once. This may be defined separately for different type of accesses and/or combination of accesses (e.g., separately defined for read accesses and write accesses).

[0052] Control circuit **204** may include a state machine or several state machines. In some embodiments, control circuit **204** may include a general purpose processor or micro-controller that is programmed to control storage module **120**. For example, a processor programmed with firmware may implement one or more state machines that govern the operation of storage module **120**. Firmware or other software for programming control circuit **204** may be stored in dedicated storage or in a reserved storage part on mass storage **202**. As another alternative, control circuit **204** may be implemented as a combination of a general purpose processor programmed with firmware or the like and special purpose circuitry that performs specific functions.

[0053] Among the aspects of storage module **120** that control circuit **204** controls is the operation of clock generation circuit **206**. In particular, using information stored in registers **212** and state information, which, in some examples, may also be stored in registers **212** or alternatively in RAM **214**, control circuit **204** supplies control information (e.g., control bits) to clock generation circuit that controls the operation of the internal clock signal.

[0054] Other functions of control circuit **204** include receiving command signals from host **100** to perform certain functions. For example, control circuit **204** may receive

command signals from host **100** to read information from, or write information to, registers **212**. For instance, control circuit **204** may receive a command to read registers **212** in a location that stores a state of storage module **120** (e.g., a power state, a programming state, etc.).

[0055] It should be understood that the architecture of FIG. 2 is an example for case of discussion only and should not be considered limiting on the disclosure. Circuits, buses, lines, modules and the like may have been simplified, left out, or otherwise combined with other components in FIG. 2. For example, storage module **120** is shown to have buses, such as internal clock line **207**, bus **205**, bus **213**, bus **211**, and bus **222**; these buses may be removed, combined, rerouted, and/or added to without departing from the embodiments described herein. As another example, the functionality of control circuit **204** may be greatly expanded over what is described above and/or the functions described above by control circuit **204** may be spread across different circuits.

[0056] FIG. 3 depicts additional example details of mass storage **202** and its connection to storage controller **200** through bus **222**. In particular, mass storage **202** may be made up of a set of memories, e.g., chips **301-304** (e.g., NAND memory chips). Each memory chip may be further divided into a set of portions. For example, memory chip **302** may be organized into two portions, e.g., plane **306** and plane **308**. Similarly, memory chip **304** may be organized into two portions, e.g., plane **310** and **312**. While not shown in FIG. 3, memory chips **301** and **303** may be similarly organized. Each portion (e.g., plane) of memory chips **301-304** may contain a certain number of blocks. For example, plane **306** may contain the odd numbered blocks of memory chip **302** and plane **308** may contain the even numbered blocks of memory chip **302**. Each block then may contain a number of pages of memory. Each page may contain a number of bits of memory. The two-plane architecture, e.g., for a memory chip, allows for simultaneous reads from, and/or writes to, two pages or the concurrent erasing of two blocks. While blocks and pages are being used to describe the physical organization of memory chips **301-304**, these terms may also be used to describe logical organization of memory in other contexts.

[0057] Bus **222** may include a set of channels, e.g., channel **314** and channel **316**. Each channel provides a connection between storage controller **200** and a memory (e.g., a memory chip or set of chips) of mass storage **202**. In particular, channel **314** provides communication (e.g., data transmission) between storage controller **200** and memory chips **301** and **302**. Channel **316** provides the communication between storage controller **200** and memory chips **303** and **304**. The example dual-channel architecture of mass storage **202** allows for higher throughput between storage controller **200** and mass storage **202**.

[0058] Other variations are possible of the hardware layout of mass storage **202** depicted in FIG. 3. For example, FIG. 4 depicts a mass storage **202** that has two memory chips **301** and **302** (not four as in FIG. 3), where memory chip **302** may be organized into two portions, e.g., plane **306** and plane **308** as previously described. Also, mass storage **202** connects to storage controller **200** through bus **222** which has one channel **314** (not two as in FIG. 3). The architectures shown in FIGS. 3 and 4 are for example purposes. The number of channels, number of memory chips in the mass storage **202**, the number of memory chips connected to each

channel, the number of planes per memory chip, etc., may vary without departing from embodiments of the present disclosure.

[0059] Other examples of possible hardware layouts of mass storage **202** may include: a set of two memory chips connected to the storage controller **200** by one channel and each memory chip having one plane; a set of four memory chips connected to the storage controller **200** by a set of two channels (e.g., each channel connecting two memory chips to the storage controller **200**) and each memory chip have a set of four planes; and one memory chip connected to the storage controller **200** by one channel and the memory chip having a set of two planes.

[0060] FIG. 5 depicts exemplary process **500** that implements a first embodiment of the present disclosure. Storage module **120** may implement the storage module portion of process **500**, which itself is an embodiment of the present disclosure.

[0061] In step **502**, host **100** sends data for storage and an identification or identifier (e.g., ID) associated with the data. The ID of the data may be part of a write command containing the data or may be an argument of a write command. Alternatively, the ID of the data may be sent, e.g., via a separate command, prior to a write command but is associated with the write command because the write command is provided subsequent to the ID of the data.

[0062] The ID of the data may be separate from the data and addresses of the data in that the ID may not be extracted from the data that is being stored nor extracted from the address(es) of the storage location(s) of the data. For example, the ID of the data may take the form of existing metadata tags in the eMMC or UFS standards. Specifically, the ID of the data may take the form of ContextIDs used in the eMMC standard or DataTAGs used in the UFS standard.

[0063] In step **504**, storage module **120** receives the data and the ID of the data. As discussed with respect to step **502**, this may happen with one command or multiple commands. For example, storage module **120** may first receive the ID of the data as part of one command. Then all subsequent data received as part of one or more write commands will be associated with that ID. As an alternative example, host **100** may send the ID of the data as part of the write command or an argument of the write command for the data. As yet another alternative example, storage module **120** may receive the data and store the data temporarily. Storage module **120** may then receive the ID of the data while the data is being temporarily store in a temporary storage location.

[0064] In step **506**, storage module **120** determines the ID of the data received in step **504**. For example, control circuit **204** may extract the ID of the data from a write command associated with the data or read the most recently sent ID of the data from a register or some other storage location.

[0065] In step **508**, storage module **120** determines which hardware resources to use to store the data based on the ID of the data. For example, storage controller **200** may use the ID of the data to select one or more of a specific memory chip, a specific channel, and/or a specific plane to use when programming the data to mass storage **202**. In cases where storage module **120** contains one type of resource, storage controller **200** may just use that single resource by default. For example, if storage controller **120** has one channel, then that one channel may always be used by default.

[0066] In one example, the ID of the data maybe a reserved contextID that is reserved to indicate that the data being written to the storage module **120** will be accessed as one of the parallel data streams that storage module **120** supports. Specifically and further to this example, if chunk1 of data is marked with an ID of “01h” and chunk2 of data is marked with an ID of “02h”, then storage module **120** may store chunk1 and chunk2 in such a manner that parallel read access to chunk1 and chunk2 is possible. As examples, storage module **120** may ensure that chunk1 is stored using a different channel than a channel used to store chunk2, that chunk1 is stored on a different memory chip than a memory chip used to store chunk2, or that chunk1 is stored in a different plane than a plane used to store chunk2. If parallel read access is possible, then command queuing of read commands to chunk1 and chunk2 may improve performance over sequential execution of the commands.

[0067] In some embodiments, host **100** may be aware of a number of possible parallel read accesses to implement the above example. Host **100** may learn this information by, for example, reading data from registers **212** as described above. Host **100** may set the ID of data according to any number of factors, such as an application associated with the data, a data type (e.g., video, images, email, database, etc.), hardware (e.g., sensors, camera, GPS, etc.) associated with the data, etc.

[0068] In another embodiment, instead of setting IDs for particular data, entire partitions or other logical units (LU) of storage module **120** may be assigned IDs. In this embodiment, storage module **120** may use the ID of the partition or LU to determine which hardware resources to use when storing data to that partition or LU.

[0069] Embodiments of the present disclosure include a storage device comprising means for receiving data to be stored in the storage device, such as for example data in line **216B** and data terminal **216A**. The storage device comprises a set of memories including at least one memory, wherein each of the memories in the set of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion. The storage device comprises a set of channels connecting the set of memories to means for controlling the storage device, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects the means for controlling the storage device to at least one of the memories in the set of memories. The means for controlling the storage device may comprise, for example, storage controller **200**. The means for receiving further receives an ID associated with the data to be stored in the storage device. The means for controlling selects, based at least in part on the ID associated with the data, at least one of (1) a memory area from the set of memories or (2) a channel from the set of channels and the means for controlling utilizes the at least one of the selected memory area or the selected channel to store the data.

[0070] Although a feature may appear to be described in connection with a particular embodiment, one skilled in the art would recognize that various features of the described embodiments may be combined. Moreover, aspects described in connection with an embodiment may stand alone.

What is claimed is:

1. A storage device comprising:

a set of memories including at least one memory, wherein each of the memories in the set of memories includes a plurality of memory circuits for storing data and is divided into a set of portions, each set of portions including at least one portion;

a controller configured to transfer the data between the set of memories and a host through an interface; and

a set of channels connecting the set of memories to the controller, wherein the set of channels includes at least one channel and each of the channels in the set of channels connects the controller to at least one of the memories in the set of memories;

wherein, the controller is further configured to select at least one of (1) a memory area from the set of memories or (2) a channel from the set of channels, based at least in part upon an ID associated with the data, and wherein, the at least one of the selected memory area or the selected channel are used to store the data.

* * * * *