

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication	20250258619
Kind Code	A1
Publication Date	August 14, 2025
Inventor(s)	Eliash; Tomer

INCOMPLETE SUPERBLOCK MANAGEMENT FOR MEMORY SYSTEMS

Abstract

Aspects of the present disclosure configure a system component, such as a memory sub-system controller. to provide superblock management based on memory component reliabilities.

Inventors:	Eliash; Tomer (Sunnyvale, CA)
Applicant:	Micron Technology, Inc. (Boise, ID)
Family ID:	89665352
Appl. No.:	19/193063
Filed:	April 29, 2025

Related U.S. Application Data

parent US continuation 18587442 20240226 parent-grant-document US 12314590 child US 19193063
parent US continuation 17875521 20220728 parent-grant-document US 11941276 child US 18587442

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

CPC G06F3/064 (20130101); G06F3/0604 (20130101); G06F3/0629 (20130101);
G06F3/0679 (20130101);

Background/Summary

PRIORITY APPLICATION [0001] This application is a continuation of U.S. application Ser. No. 18/587,442, filed Feb. 26, 2024, which is a continuation of U.S. application Ser. No. 17/875,521, filed Jul. 28, 2022, all of which are incorporated herein by reference in their entirety.

TECHNICAL FIELD

[0002] Embodiments of the disclosure relate generally to memory sub-systems and, more specifically, to providing adaptive media management for memory components, such as memory dies.

BACKGROUND

[0003] A memory sub-system can be a storage system, such as a solid-state drive (SSD), and can include one or more memory components that store data. The memory components can be, for example, non-volatile memory components and volatile memory components. In general, a host system can utilize a memory sub-system to store data on the memory components and to retrieve data from the memory components.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] The present disclosure will be understood more fully from the detailed description given below and from the accompanying drawings of various embodiments of the disclosure.

[0005] FIG. 1 is a block diagram illustrating an example computing environment including a memory sub-system, in accordance with some embodiments of the present disclosure.

[0006] FIG. 2 is a block diagram of an example media operations manager, in accordance with some implementations of the present disclosure.

[0007] FIG. 3 is a block diagram of an example plurality of superblocks, in accordance with some implementations of the present disclosure.

[0008] FIG. 4 is a flow diagram of an example method to perform memory operations on a plurality of superblocks before an incomplete superblock, in accordance with some implementations of the present disclosure.

[0009] FIG. 5 provides an interaction diagram illustrating interactions between components of the computing environment in the context of some embodiments in which a media operations manager is used.

[0010] FIG. 6 is a block diagram illustrating a diagrammatic representation of a machine in the form of a computer system within which a set of instructions can be executed for causing the machine to perform any one or more of the methodologies discussed herein, in accordance with some embodiments of the present disclosure.

DETAILED DESCRIPTION

[0011] Aspects of the present disclosure configure a system component, such as a memory sub-system controller, to perform memory operations on complete superblocks (CS) before incomplete superblocks (ICS). The memory sub-system controller can access configuration data associated with a set of memory components. The configuration data can include a table that associates individual groups of memory components with respective reliability grades that describe a data retention parameter, a read disturb parameter, an error rate, a leakage current, a cross temperature parameter, and/or an endurance parameter of the individual group of memory components. Based on configuration data (which can be stored on the memory components and/or a memory of the controller), the controller can group blocks of the set of memory components into complete superblocks and incomplete superblocks. Specifically, complete superblocks can include a set of

blocks that each have a reliability grade above a threshold, and incomplete superblocks can include a first portion of blocks with reliability grades above the threshold and a second portion of blocks with reliability grades below a threshold. The controller can allocate (e.g., perform one or more memory operations on) the complete superblocks before the incomplete superblocks. Namely, once all of the complete superblocks have been allocated, then the controller can start allocating the incomplete superblocks. This ensures that performance of the memory system remains optimal by only allocating and performing memory operations on incomplete superblocks after all of the complete superblocks have been allocated. This improves the overall efficiency of operating the memory sub-system.

[0012] A memory sub-system can be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of storage devices and memory modules are described below in conjunction with FIG. 1. In general, a host system can utilize a memory sub-system that includes one or more memory components, such as memory devices (e.g., memory dies) that store data. The host system can send access requests (e.g., write command, read command) to the memory sub-system, such as to store data at the memory sub-system and to read data from the memory sub-system. The data (or set of data) specified by the host is hereinafter referred to as “host data,” “application data,” or “user data.”

[0013] The memory sub-system can initiate media management operations, such as a write operation, on host data that is stored on a memory device. For example, firmware of the memory sub-system may rewrite previously written host data from a location on a memory device to a new location as part of garbage collection management operations. The data that is rewritten, for example as initiated by the firmware, is hereinafter referred to as “garbage collection data.” “User data” can include host data and garbage collection data. “System data” hereinafter refers to data that is created and/or maintained by the memory sub-system for performing operations in response to host requests and for media management. Examples of system data include, and are not limited to, system tables (e.g., logical-to-physical address mapping table), data from logging, scratch pad data, and so forth.

[0014] Many different media management operations can be performed on a memory device. For example, the media management operations can include different scan rates, different scan frequencies, different wear leveling, different read disturb management, different near miss error correction (ECC), and/or different dynamic data refresh. Wear leveling ensures that all blocks in a memory component approach their defined erase-cycle budget at the same time, rather than some blocks approaching it earlier. Read disturb management counts all of the read operations to the memory component. If a certain threshold is reached, the surrounding regions are refreshed. Near-miss ECC refreshes all data read by the application that exceeds a configured threshold of errors. Dynamic data-refresh scan reads all data and identifies the error status of all blocks as a background operation. If a certain threshold of errors per block or ECC unit is exceeded in this scan-read, a refresh operation is triggered.

[0015] A memory device can be a non-volatile memory device. A non-volatile memory device is a package of one or more dice (or dies). Each die can be comprised of one or more planes. For some types of non-volatile memory devices (e.g., negative- and (NAND) devices), each plane is comprised of a set of physical blocks. For some memory devices, blocks are the smallest area than can be erased. Each block is comprised of a set of pages. Each page is comprised of a set of memory cells, which store bits of data. The memory devices can be raw memory devices (e.g., NAND), which are managed externally, for example, by an external controller. The memory devices can be managed memory devices (e.g., managed NAND), which is a raw memory device combined with a local embedded controller for memory management within the same memory device package.

[0016] There are challenges in efficiently managing or performing media management operations on typical memory devices. Specifically, certain memory devices, such as NAND flash devices,

include large die-by-die reliability variation. As the technology for such memory devices continues to be scaled down, this die-by-die reliability variation becomes more pronounced and problematic in performing memory management. Current memory systems (e.g., SSD drive or die package systems) associate all of the memory devices in the memory system with a certain reliability specification. In some cases, each block of each memory device is associated with a reliability grade or specification which is used to determine whether the block is a good block or a bad block. Good blocks are those that have reliability grades above a threshold and bad blocks are blocks that have reliability grades below a threshold. The reliability grades can be set at manufacture or during operation of the memory devices, such as by measuring the data retention and/or error rate associated with particular blocks.

[0017] Typical memory systems leverage superblocks, which are a collection of blocks across memory planes and/or dies. Namely, each superblock can be of equal size and can include a respective collection of blocks across multiple planes and/or dies. The superblocks, when allocated, allow a controller to simultaneously write data to a large portion of memory spanning multiple blocks (across multiple planes and/or dies) with a single address. Sometimes, superblocks include bad blocks or blocks that have reliability grades that are below a threshold. These can be referred to as incomplete superblocks. Typical systems allocate these incomplete superblocks in the same manner as complete superblocks (e.g., superblocks that include only good blocks that have reliability grades above the threshold). This usually results in poor memory performance as performing memory operations on incomplete superblocks can result in a greater quantity of errors or unreliable memory behavior. Also, some of the bad blocks in the incomplete superblocks cannot be used to perform memory operations. This further reduces the efficiency of allocating such incomplete superblocks because less memory space is available for performing memory operations (e.g., the memory operations can only be performed on the good blocks of the incomplete superblocks). As such, applying a one-size-fits-all approach to memory systems that have a mix of complete and incomplete superblocks is inefficient and results in poor or unreliable memory performance.

[0018] Aspects of the present disclosure address the above and other deficiencies by providing a memory controller that can allocate all of the complete superblocks before allocating incomplete superblocks to perform memory management operations on the respective memory components. For example, the memory controller can group a plurality of sets of blocks of the set of memory components into respective ones of a plurality of superblocks and determine that one or more blocks of an individual superblock of the plurality of superblocks are associated with a reliability grade that is lower than a threshold. In such cases, the memory controller designates the individual superblock as an incomplete superblock and performs one or more memory operations on a portion of the plurality of superblocks including complete superblocks before the incomplete superblock. The one or more memory operations can include at least one of a wear leveling operation, a data read or write operation, and/or a garbage collection operation. As a result, the memory controller can tailor the memory management operations to their particular reliability grades of the memory components without having to sacrifice performance by waiting to allocate incomplete superblocks until that additional space is needed (e.g., all of the complete superblocks have been allocated). This increases the efficiency of operating memory systems.

[0019] In some examples, the memory controller accesses configuration data. The configuration data includes a table that associates individual blocks of the set of memory components with respective reliability grades. The reliability grade describes at least one of a data retention parameter, a read disturb parameter, an error rate, a leakage current, a cross temperature parameter, and/or an endurance parameter. The memory controller can designate the incomplete superblocks based on the configuration data.

[0020] In some examples, the memory controller allocates each of the complete superblocks for storing data before allocating any blocks of the incomplete superblock. A first set of the plurality of

sets of blocks of a first superblock can be distributed across multiple memory dies or memory planes. The memory controller can store a table that includes a first entry associating a first set of the plurality of sets of blocks with a first superblock address. The table can include a second entry associating a second set of the plurality of sets of blocks with a second superblock address. A total storage space can correspond to the plurality of superblocks including the incomplete superblock corresponding to an exported capacity of the memory sub-system.

[0021] In some examples, the memory controller can receive a request to perform a memory operation. In response, the memory controller determines that each of the plurality of superblocks excluding the incomplete superblock has been previously allocated and performs the memory operation on a set of blocks of the incomplete superblock. The memory controller allocates the incomplete superblock to the memory operation. In some cases, the incomplete superblock is a first incomplete superblock. In such cases, the memory controller determines that a first collection of blocks of a second superblock of the plurality of superblocks are associated with an individual reliability grade that is lower than the threshold, the second superblock including a second collection of blocks associated with individual reliability grades that transgress the threshold. The memory controller designates the second superblock as a second incomplete superblock.

[0022] In some examples, the memory controller determines a quantity of blocks of the first incomplete superblock that are associated with the reliability grade that is lower than the threshold and selects a group of the second collection of the blocks of the second incomplete superblock that corresponds to the quantity of blocks. For example, if the first incomplete superblock includes two bad blocks, the memory controller selects a group of two blocks from the good blocks of the second incomplete superblock. Then, the memory controller associates the selected group of the second collection of blocks with the first incomplete superblock. By doing so, the memory controller replaces the quantity of blocks of the first incomplete superblock with the selected group of the second collection of blocks. In some examples, the memory controller performs the memory operation on the first incomplete superblock in multiple passes. Namely, a first address of the first incomplete superblock is used in a first of the multiple passes and a second address of the second incomplete superblock is used in a second of the multiple passes.

[0023] Though various embodiments are described herein as being implemented with respect to a memory sub-system (e.g., a controller of the memory sub-system), some or all of the portions of an embodiment can be implemented with respect to a host system, such as a software application or an operating system of the host system.

[0024] FIG. 1 illustrates an example computing environment **100** including a memory sub-system **110**, in accordance with some examples of the present disclosure. The memory sub-system **110** can include media, such as memory components **112A** to **112N** (also hereinafter referred to as “memory devices”). The memory components **112A** to **112N** can be volatile memory devices, non-volatile memory devices, or a combination of such. The memory components **112A** to **112N** can be implemented by individual dies, such that a first memory component **112A** can be implemented by a first memory die (or a first collection of memory dies) and a second memory component **112N** can be implemented by a second memory die (or a second collection of memory dies).

[0025] In some examples, the first memory component **112A**, block or page of the first memory component **112A**, or group of memory components including the first memory component **112A** can be associated with a first reliability (capability) grade, value, or measure. The terms “reliability grade,” “value,” and “measure” are used interchangeably throughout and can have the same meaning. The second memory component **112N** or group of memory components including the second memory component **112N** can be associated with a second reliability (capability) grade, value, or measure. In some examples, each memory component **112A** to **112N** can store respective configuration data that specifies the respective reliability grade. In some examples, a memory or register can be associated with all of the memory components **112A** to **112N** that can store a table that maps different groups, bins, or sets of the memory components **112A** to **112N** to respective

reliability grades.

[0026] In some embodiments, a block within the first memory component **112A** can be grouped with a block within the second memory component **112N** to form a superblock. Superblocks can be addressed collectively using a single address. In such cases, a table can store the association between the single address and each of the blocks of the first memory component **112A** and second memory component **112N** associated with that single address. In some embodiments, some of the blocks of the superblock can have reliability grades that are below a threshold. These can be referred to as bad blocks. When a superblock includes one or more bad blocks, the superblock is referred to as an incomplete superblock. A superblock that includes no bad blocks and only includes good blocks is referred to as a complete superblock.

[0027] In some embodiments, the memory sub-system **110** is a storage system. A memory sub-system **110** can be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of a storage device include a SSD, a flash drive, a universal serial bus (USB) flash drive, an embedded Multi-Media Controller (eMMC) drive, a Universal Flash Storage (UFS) drive, and a hard disk drive (HDD). Examples of memory modules include a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), and a non-volatile dual in-line memory module (NVDIMM).

[0028] The computing environment **100** can include a host system **120** that is coupled to a memory system. The memory system can include one or more memory sub-systems **110**. In some embodiments, the host system **120** is coupled to different types of memory sub-system **110**. FIG. 1 illustrates one example of a host system **120** coupled to one memory sub-system **110**. The host system **120** uses the memory sub-system **110**, for example, to write data to the memory sub-system **110** and read data from the memory sub-system **110**. As used herein, “coupled to” generally refers to a connection between components, which can be an indirect communicative connection or direct communicative connection (e.g., without intervening components), whether wired or wireless, including connections such as electrical, optical, magnetic, and so forth.

[0029] The host system **120** can be a computing device such as a desktop computer, laptop computer, network server, mobile device, embedded computer (e.g., one included in a vehicle, industrial equipment, or a networked commercial device), or such computing device that includes a memory and a processing device. The host system **120** can include or be coupled to the memory sub-system **110** so that the host system **120** can read data from or write data to the memory sub-system **110**. The host system **120** can be coupled to the memory sub-system **110** via a physical host interface. Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, a USB interface, a Fibre Channel interface, a Serial Attached SCSI (SAS) interface, and the like. The physical host interface can be used to transmit data between the host system **120** and the memory sub-system **110**. The host system **120** can further utilize an NVM Express (NVMe) interface to access the memory components **112A** to **112N** when the memory sub-system **110** is coupled with the host system **120** by the PCIe interface. The physical host interface can provide an interface for passing control, address, data, and other signals between the memory sub-system **110** and the host system **120**.

[0030] The memory components **112A** to **112N** can include any combination of the different types of non-volatile memory components and/or volatile memory components. An example of non-volatile memory components includes a NAND-type flash memory. Each of the memory components **112A** to **112N** can include one or more arrays of memory cells such as single-level cells (SLCs) or multi-level cells (MLCs) (e.g., TLCs or QLCs). In some embodiments, a particular memory component **112** can include both an SLC portion and an MLC portion of memory cells. Each of the memory cells can store one or more bits of data (e.g., blocks) used by the host system **120**. Although non-volatile memory components such as NAND-type flash memory are described, the memory components **112A** to **112N** can be based on any other type of memory, such as a

volatile memory. In some embodiments, the memory components **112A** to **112N** can be, but are not limited to, random access memory (RAM), read-only memory (ROM), dynamic random access memory (DRAM), synchronous dynamic random access memory (SDRAM), phase change memory (PCM), magnetoresistive random access memory (MRAM), negative-or (NOR) flash memory, electrically erasable programmable read-only memory (EEPROM), and a cross-point array of non-volatile memory cells. A cross-point array of non-volatile memory cells can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many flash-based memories, cross-point non-volatile memory can perform a write-in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased. Furthermore, the memory cells of the memory components **112A** to **112N** can be grouped as memory pages or blocks that can refer to a unit of the memory component **112** used to store data. For example, a single first row that spans memory components **112A** to **112N** can correspond to or be grouped as a first superblock and a single second row that spans memory components **112A** to **112N** can correspond to or be grouped as a second superblock. If the single first row includes all good blocks (e.g., each block in the single first row has a reliability grade above a threshold), the first superblock is a first complete superblock. If the single first row includes some bad blocks (e.g., one or more blocks in the single first row have a reliability grade below a threshold), the first superblock is a first incomplete superblock.

[0031] A memory sub-system controller **115** can communicate with the memory components **112A** to **112N** to perform memory operations such as reading data, writing data, or erasing data at the memory components **112A** to **112N** and other such operations. The memory sub-system controller **115** can communicate with the memory components **112A** to **112N** to perform various memory management operations, such as different scan rates, different scan frequencies, different wear leveling, different read disturb management, different near miss ECC operations, and/or different dynamic data refresh.

[0032] The memory sub-system controller **115** can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, or a combination thereof. The memory sub-system controller **115** can be a microcontroller, special-purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), etc.), or another suitable processor. The memory sub-system controller **115** can include a processor (processing device) **117** configured to execute instructions stored in local memory **119**. In the illustrated example, the local memory **119** of the memory sub-system controller **115** includes an embedded memory configured to store instructions for performing various processes, operations, logic flows, and routines that control operation of the memory sub-system **110**, including handling communications between the memory sub-system **110** and the host system **120**. In some embodiments, the local memory **119** can include memory registers storing memory pointers, fetched data, and so forth. The local memory **119** can also include ROM for storing microcode. While the example memory sub-system **110** in FIG. 1 has been illustrated as including the memory sub-system controller **115**, in another embodiment of the present disclosure, a memory sub-system **110** may not include a memory sub-system controller **115**, and can instead rely upon external control (e.g., provided by an external host, or by a processor **117** or controller separate from the memory sub-system **110**).

[0033] In general, the memory sub-system controller **115** can receive commands or operations from the host system **120** and can convert the commands or operations into instructions or appropriate commands to achieve the desired access to the memory components **112A** to **112N**. In some examples, the commands or operations received from the host system **120** can specify configuration data for the memory components **112A** to **112N**. The configuration data can describe the reliability grades associated with different groups of the memory components **112A** to **112N** and/or different blocks within each of the memory components **112A** to **112N**. In some cases, the

reliability grades are dynamic and can be updated by the memory sub-system controller **115** in response to determining that certain error rates are reached that transgress an error rate threshold. For example, a good block can become a bad block if that good block starts having error rates that transgress the threshold. In such cases, the configuration data is updated and any superblock that includes that now bad block is designated as an incomplete superblock.

[0034] The memory sub-system controller **115** can be responsible for other memory management operations, such as wear leveling operations, garbage collection operations, error detection and ECC operations, encryption operations, caching operations, and address translations. The memory sub-system controller **115** can further include host interface circuitry to communicate with the host system **120** via the physical host interface. The host interface circuitry can convert the commands received from the host system **120** into command instructions to access the memory components **112A** to **112N** as well as convert responses associated with the memory components **112A** to **112N** into information for the host system **120**.

[0035] The memory sub-system **110** can also include additional circuitry or components that are not illustrated. In some embodiments, the memory sub-system **110** can include a cache or buffer (e.g., DRAM or other temporary storage location or device) and address circuitry (e.g., a row decoder and a column decoder) that can receive an address from the memory sub-system controller **115** and decode the address to access the memory components **112A** to **112N**.

[0036] The memory devices can be raw memory devices (e.g., NAND), which are managed externally, for example, by an external controller (e.g., memory sub-system controller **115**). The memory devices can be managed memory devices (e.g., managed NAND), which is a raw memory device combined with a local embedded controller (e.g., local media controllers) for memory management within the same memory device package. Any one of the memory components **112A** to **112N** can include a media controller (e.g., media controller **113A** and media controller **113N**) to manage the memory cells of the memory component (e.g., to perform one or more memory management operations), to communicate with the memory sub-system controller **115**, and to execute memory requests (e.g., read or write) received from the memory sub-system controller **115**.

[0037] The memory sub-system controller **115** can include a media operations manager **122**. The media operations manager **122** can be configured to allocate all of the complete superblocks (superblocks without any bad blocks) before allocating incomplete superblocks (superblocks with at least one bad block) to perform memory operations on the respective memory components **112A** to **112N**. For example, the memory controller can group a plurality of sets of blocks of the set of memory components **112A** to **112N** into respective ones of a plurality of superblocks and determine that one or more blocks of an individual superblock of the plurality of superblocks are associated with a reliability grade that is lower than a threshold. In such cases, the memory controller designates the individual superblock as an incomplete superblock and performs one or more memory operations on the complete superblocks before the incomplete superblock. The one or more memory operations can include at least one of a wear leveling operation, a data read or write operation, and/or garbage collection operation. As a result, the memory controller can tailor the memory management operations to their particular reliability grades of the memory components without having to sacrifice performance by waiting to allocate incomplete superblocks until that additional space is needed (e.g., all of the complete superblocks have been allocated). This increases the efficiency of operating memory systems.

[0038] As an example, the media operations manager **122** can receive a request to perform one or more memory operations on the memory components **112A** to **112N**. In response, the media operations manager **122** can access a table of allocated superblocks to determine whether there exist any complete superblocks that have not been allocated. In response to identifying a complete superblock that has not yet been allocated, the media operations manager **122** allocates and performs the one or more memory operations on the identified complete superblock. In response to determining that all of the complete superblocks have been allocated, the media operations

manager **122** selects one or more incomplete superblocks to allocate and perform the one or more memory operations. In some cases, the media operations manager **122** combines multiple incomplete superblocks to generate a pair of incomplete superblocks that together include the same quantity of good blocks as one of the complete superblocks. This increases the overall efficiency of operating the memory device.

[0039] Depending on the embodiment, the media operations manager **122** can comprise logic (e.g., a set of transitory or non-transitory machine instructions, such as firmware) or one or more components that causes the media operations manager **122** to perform operations described herein. The media operations manager **122** can comprise a tangible or non-tangible unit capable of performing operations described herein. Further details with regards to the operations of the media operations manager **122** are described below.

[0040] FIG. 2 is a block diagram of an example media operations manager **200**, in accordance with some implementations of the present disclosure. As illustrated, the media operations manager **122** includes configuration data **220**, a superblock designation table **230**, and a superblock allocation module **240**. For some embodiments, the media operations manager **122** can differ in components or arrangement (e.g., less or more components) from what is illustrated in FIG. 2.

[0041] The configuration data **220** accesses and/or stores configuration data associated with the memory components **112A** to **112N**. In some examples, the configuration data **220** is programmed into the media operations manager **122**. For example, the media operations manager **122** can communicate with the memory components **112A** to **112N** to obtain the configuration data and store the configuration data **220** locally on the media operations manager **122**. In some examples, the media operations manager **122** communicates with the host system **120**. The host system **120** receives input from an operator or user that specifies parameters including reliability grades of different bins, groups, blocks, or sets of the memory components **112A** to **112N**. The media operations manager **122** receives configuration data from the host system **120** and stores the configuration data in the configuration data **220**.

[0042] In some examples, the media operations manager **122** performs one or more test operations on different groups or blocks of the memory components **112A** to **112N**. The test operations are configured to determine reliability of each block of the memory components **112A** to **112N**. Based on a result of the test operations, the media operations manager **122** can store or update the reliability grades stored in the configuration data **220** for the different groups or blocks of memory components **112A** to **112N**. In some examples, the media operations manager **122** can periodically or routinely perform the test operations. The media operations manager **122** can determine that an individual memory component **112A** is associated with a first reliability grade based on the configuration data **220**. The media operations manager **122** can perform a set of test operations on the individual memory component **112A** and can determine, based on a result of the test operations, that the reliability grade of the individual memory component **112A** has increased or decreased and is now a second reliability grade. The media operations manager **122** can, in response, transition the individual memory components **112A** from being associated with a first group of individual memory components **112A** to **112N** to a second group of individual memory components **112A** to **112N** that is associated with the second reliability grade.

[0043] In some examples, the media operations manager **122** processes the configuration data **220** to generate a superblock designation table **230** that lists addresses of complete superblocks and addresses of incomplete superblocks. For example, a first address in the table can be associated with a first row of good blocks across multiple dies or planes representing a first complete superblock. A second address in the table can be associated with a second row of good blocks across multiple dies or planes representing a second complete superblock. Specifically, as shown in FIG. 3, a configuration **300** of the memory components **112A** to **112N** is provided. The configuration includes a first die **310** (e.g., the first memory components **112A**) and a second die **320** (e.g., the second memory components **112N**).

[0044] The media operations manager **122** processes the configuration data **220** to determine that a first row **316** of blocks across the first die **310** and the second die **320** have a reliability grade that corresponds to a threshold (e.g., have error rates that are below a specified error rate threshold). In such cases, the media operations manager **122** stores a first complete superblock identifier **312** in the superblock designation table **230**. The first complete superblock identifier **312** can include an address that accesses the entire first row **316** of blocks. Similarly, the media operations manager **122** associates a second complete superblock identifier **314** associated with another entire row of good blocks.

[0045] In some cases, the media operations manager **122** identifies a bad block **334** in an individual row of blocks of a superblock that also includes a set of good blocks **331**. In response, the media operations manager **122** stores a first incomplete superblock identifier **330** in the superblock designation table **230**. Similarly, the media operations manager **122** identifies a bad block in another individual row of blocks of another superblock that also includes a set of good blocks **336**. In response, the media operations manager **122** stores a second incomplete superblock identifier **332** in the superblock designation table **230**.

[0046] In some examples, the media operations manager **122** receives a request to perform a memory operation on the superblocks stored in the superblock designation table **230**. In response, the superblock allocation module **240** searches the superblock designation table **230** to determine if there exist any unallocated complete superblocks. In response to identifying an individual complete superblock that has been unallocated, the superblock allocation module **240** stores an indicator in the superblock designation table **230** in association with the individual complete superblock that the individual complete superblock has been allocated. Also, the superblock allocation module **240** processes the requested memory operation on the individual complete superblock that has been allocated.

[0047] In some cases, the superblock allocation module **240** searches the superblock designation table **230** to determine that all of the complete superblocks have been allocated. In such cases, the superblock allocation module **240** identifies an incomplete superblock that has not been allocated. The superblock allocation module **240** stores an indicator in the superblock designation table **230** in association with the identified incomplete superblock that the identified incomplete superblock has been allocated. Also, the superblock allocation module **240** processes the requested memory operation on the identified incomplete superblock that has been allocated.

[0048] In some examples, the superblock allocation module **240** combines one incomplete superblock with a portion of another incomplete superblock to provide a superblock that includes the same number of good blocks as one of the previously allocated complete superblocks. To do so, the superblock allocation module **240** computes how many bad blocks are included in the identified incomplete superblock. For example, the superblock allocation module **240** can determine that the first incomplete superblock identifier **330** includes only one bad block **334**. In such cases, the superblock allocation module **240** identifies another incomplete superblock, such as the second incomplete superblock identifier **332**. The superblock allocation module **240** selects one good block **336** from the set of blocks grouped into the second incomplete superblock identifier **332**. The superblock allocation module **240** replaces the bad block **334** from the set of blocks grouped into the first incomplete superblock identifier **330** with the good block **336**. Namely, the superblock allocation module **240** stores, in the superblock designation table **230**, the list of identifiers or addresses of each good block in the row of blocks associated with the first incomplete superblock identifier **330** and the address or identifier of the selected good block **336** of the second incomplete superblock identifier **332**. When a request to perform an operation on the first incomplete superblock identifier **330** is received, the request is performed in two phases. The first phase includes first operating on all of the blocks of the row associated with the first incomplete superblock identifier **330** and the second phase includes operating on the single good block **336** of the second incomplete superblock identifier **332**.

[0049] FIG. 4 is a flow diagram of an example method **400** to perform memory operations on a plurality of superblocks before an incomplete superblock, in accordance with some implementations of the present disclosure. The method **400** can be performed by processing logic that can include hardware (e.g., a processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, an integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **400** is performed by the media operations manager **122** of FIG. 1. Although the processes are shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other process flows are possible.

[0050] Referring now FIG. 4, the method (or process) **400** begins at operation **405**, with a media operations manager **122** of a memory sub-system (e.g., memory sub-system **110**) grouping a plurality of sets of blocks of the set of memory components into respective ones of a plurality of superblocks. Then, at operation **410**, the media operations manager **122** of the memory sub-system determines that one or more blocks of an individual superblock of the plurality of superblocks are associated with a reliability grade that is lower than a threshold. Thereafter, at operation **415**, the media operations manager **122** designates the individual superblock as an incomplete superblock. The media operations manager **122**, at operation **420**, performs one or more memory operations on a portion of the plurality of superblocks comprising complete superblocks before the incomplete superblock.

[0051] FIG. 5 provides an interaction diagram illustrating interactions between components of the computing environment **100** in the context of some embodiments. The operations of the method **400** can be performed by processing logic that can include hardware (e.g., a processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, an integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method is performed by the memory sub-system controller **115** and/or by the media operations manager **122**. Although the operations are shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment.

[0052] In the context of the example illustrated in FIG. 5, the host system can comprise the host system **120**, the processor of the memory sub-system controller can comprise the processor **117** of the memory sub-system **110**, and the one or more memory components can comprise at least one of the memory components **112A** to **112N**.

[0053] As shown in FIG. 5, at operation **502**, the host system (e.g., via a processor of the host system **120**) generates a request to read/write data to the memory components and, at operation **512**, the processor of the memory sub-system controller receives the request. The processor of the memory sub-system controller generates instructions to access a superblock allocation table, such as the superblock designation table **230**, at operation **514**. The processor of the memory sub-system controller generates instructions to allocate the plurality of superblocks that are complete superblocks before allocating any of the incomplete superblocks at operation **516**.

[0054] The processor of the memory sub-system controller then performs the requested read/write data operations on the allocated superblock (e.g., the complete superblock, the incomplete superblock, and/or multiple combined incomplete superblocks). The memory components **112A** to **112N** receive the instructions and, at operation **532**, the memory components **112A** to **112N** receive and store the user data in the superblocks.

[0055] In view of the disclosure above, various examples are set forth below. It should be noted that one or more features of an example, taken in isolation or combination, should be considered within the disclosure of this application.

[0056] Example 1: a system comprising: a set of memory components of a memory sub-system; and a processing device operatively coupled to the set of memory components, the processing device being configured to perform operations comprising: grouping a plurality of sets of blocks of the set of memory components into respective ones of a plurality of superblocks; determining that one or more blocks of an individual superblock of the plurality of superblocks are associated with a reliability grade that is lower than a threshold; designating the individual superblock as an incomplete superblock; and performing one or more memory operations on a portion of the plurality of superblocks comprising complete superblocks before the incomplete superblock.

[0057] Example 2: the system of Example 1 wherein the operations comprise accessing configuration data, wherein the configuration data comprises a table that associates individual blocks of the set of memory components with respective reliability grades.

[0058] Example 3: the system of Examples 1 or 2, wherein the reliability grade describes at least one of a data retention parameter, a read disturb parameter, an error rate, a leakage current, a cross temperature parameter, or an endurance parameter.

[0059] Example 4: the system of any one of Examples 1-3, wherein the operations comprise allocating each of the complete superblocks for storing data before allocating any blocks of the incomplete superblock.

[0060] Example 5: the system of any one of Examples 1-4, wherein a first set of the plurality of sets of blocks of a first superblock is distributed across multiple memory dies.

[0061] Example 6: the system of any one of Examples 1-5, wherein a first set of the plurality of sets of blocks of a first superblock is distributed across multiple memory planes.

[0062] Example 7: the system of any one of Examples 1-6, wherein the operations comprise storing a table, the table comprising a first entry associating a first set of the plurality of sets of blocks with a first superblock address, and the table comprising a second entry associating a second set of the plurality of sets of blocks with a second superblock address.

[0063] Example 8: the system of any one of Examples 1-7, wherein a total storage space corresponding to the plurality of superblocks including the incomplete superblock corresponds to an exported capacity of the memory sub-system.

[0064] Example 9: the system of any one of Examples 1-8, wherein the one or more memory operations comprise at least one of a wear leveling operation, a data read or write operation, or a garbage collection operation.

[0065] Example 10: the system of any one of Examples 1-9, wherein the operations comprise: receiving a request to perform a memory operation; determining that each of the plurality of superblocks excluding the incomplete superblock has been previously allocated; and in response to determining that each of the plurality of superblocks excluding the incomplete superblock has been previously allocated, performing the memory operation on a set of blocks of the incomplete superblock.

[0066] Example 11: the system of Example 10, wherein the operations comprise: allocating the incomplete superblock to the memory operation.

[0067] Example 12: the system of Example 10, wherein the incomplete superblock is a first incomplete superblock, and wherein the operations comprise: determining that a first collection of blocks of a second superblock of the plurality of superblocks are associated with an individual reliability grade that is lower than the threshold, the second superblock comprising a second collection of blocks associated with individual reliability grades that transgress the threshold; and designating the second superblock as a second incomplete superblock.

[0068] Example 13: the system of Example 12, wherein the operations comprise: determining a quantity of blocks of the first incomplete superblock that are associated with the reliability grade

that is lower than the threshold; selecting a group of the second collection of the blocks of the second incomplete superblock that corresponds to the quantity of blocks; and associating the selected group of the second collection of blocks with the first incomplete superblock.

[0069] Example 14: the system of Example 13, wherein the operations comprise: replacing the quantity of blocks of the first incomplete superblock with the selected group of the second collection of blocks.

[0070] Example 15: the system of Example 14, wherein the operations comprise: performing the memory operation on the first incomplete superblock in multiple passes, wherein a first address of the first incomplete superblock is used in a first of the multiple passes, and wherein a second address of the second incomplete superblock is used in a second of the multiple passes.

[0071] Methods and computer-readable storage medium with instructions for performing any one of the above Examples.

[0072] FIG. 6 illustrates an example machine in the form of a computer system **600** within which a set of instructions can be executed for causing the machine to perform any one or more of the methodologies discussed herein. In some embodiments, the computer system **600** can correspond to a host system (e.g., the host system **120** of FIG. 1) that includes, is coupled to, or utilizes a memory sub-system (e.g., the memory sub-system **110** of FIG. 1) or can be used to perform the operations of a controller (e.g., to execute an operating system to perform operations corresponding to the media operations manager **122** of FIG. 1). In alternative embodiments, the machine can be connected (e.g., networked) to other machines in a local area network (LAN), an intranet, an extranet, and/or the Internet. The machine can operate in the capacity of a server or a client machine in a client-server network environment, as a peer machine in a peer-to-peer (or distributed) network environment, or as a server or a client machine in a cloud computing infrastructure or environment.

[0073] The machine can be a personal computer (PC), a tablet PC, a set-top box (STB), a Personal Digital Assistant (PDA), a cellular telephone, a web appliance, a server, a network router, a network switch, a network bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

[0074] The example computer system **600** includes a processing device **602**, a main memory **604** (e.g., ROM, flash memory, DRAM such as SDRAM or Rambus DRAM (RDRAM), etc.), a static memory **606** (e.g., flash memory, static random access memory (SRAM), etc.), and a data storage system **618**, which communicate with each other via a bus **630**.

[0075] The processing device **602** represents one or more general-purpose processing devices such as a microprocessor, a central processing unit, or the like. More particularly, the processing device **602** can be a complex instruction set computing (CISC) microprocessor, a reduced instruction set computing (RISC) microprocessor, a very long instruction word (VLIW) microprocessor, a processor implementing other instruction sets, or processors implementing a combination of instruction sets. The processing device **602** can also be one or more special-purpose processing devices such as an ASIC, a FPGA, a digital signal processor (DSP), a network processor, or the like. The processing device **602** is configured to execute instructions **626** for performing the operations and steps discussed herein. The computer system **600** can further include a network interface device **608** to communicate over a network **620**.

[0076] The data storage system **618** can include a machine-readable storage medium **624** (also known as a computer-readable medium) on which is stored one or more sets of instructions **626** or software embodying any one or more of the methodologies or functions described herein. The instructions **626** can also reside, completely or at least partially, within the main memory **604** and/or within the processing device **602** during execution thereof by the computer system **600**, the

main memory **604** and the processing device **602** also constituting machine-readable storage media. The machine-readable storage medium **624**, data storage system **618**, and/or main memory **604** can correspond to the memory sub-system **110** of FIG. **1**.

[0077] In one embodiment, the instructions **626** implement functionality corresponding to the media operations manager **122** of FIG. **1**. While the machine-readable storage medium **624** is shown in an example embodiment to be a single medium, the term “machine-readable storage medium” should be taken to include a single medium or multiple media that store the one or more sets of instructions. The term “machine-readable storage medium” shall also be taken to include any medium that is capable of storing or encoding a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure. The term “machine-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media, and magnetic media.

[0078] Some portions of the preceding detailed descriptions have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the ways used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

[0079] It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. The present disclosure can refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system's memories or registers or other such information storage systems.

[0080] The present disclosure also relates to an apparatus for performing the operations herein. This apparatus can be specially constructed for the intended purposes, or it can include a general-purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program can be stored in a computer-readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic-optical disks; ROMs; RAMs; erasable programmable read-only memories (EPROMs); EEPROMs; magnetic or optical cards; or any type of media suitable for storing electronic instructions, each coupled to a computer system bus.

[0081] The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general-purpose systems can be used with programs in accordance with the teachings herein, or it can prove convenient to construct a more specialized apparatus to perform the method. The structure for a variety of these systems will appear as set forth in the description above. In addition, the present disclosure is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages can be used to implement the teachings of the disclosure as described herein.

[0082] The present disclosure can be provided as a computer program product, or software, that can include a machine-readable medium having stored thereon instructions, which can be used to program a computer system (or other electronic devices) to perform a process according to the present disclosure. A machine-readable medium includes any mechanism for storing information in a form readable by a machine (e.g., a computer). In some embodiments, a machine-readable (e.g., computer-readable) medium includes a machine-readable (e.g., computer-readable) storage

medium such as a ROM, RAM, magnetic disk storage media, optical storage media, flash memory components, and so forth.

[0083] In the foregoing specification, embodiments of the disclosure have been described with reference to specific example embodiments thereof. It will be evident that various modifications can be made thereto without departing from the broader spirit and scope of embodiments of the disclosure as set forth in the following claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

Claims

1. A memory sub-system comprising: a set of memory components; and a processing device operatively coupled to the set of memory components, the processing device programmed to perform operations comprising: designating an individual portion of the set of memory components as an incomplete portion based on a result of determining whether one or more sub-portions of the individual portion satisfy a reliability criterion; and performing one or more memory operations on an individual region of one or more complete portions of the set of memory components before the incomplete portion.
2. The memory sub-system of claim 1, the operations comprising: grouping a plurality of sets of blocks of the set of memory components into respective ones of a plurality of superblocks, the reliability criterion comprising a reliability grade that is lower than a threshold.
3. The memory sub-system of claim 2, wherein the reliability grade describes at least one of a data retention parameter, a read disturb parameter, an error rate, a leakage current, a cross temperature parameter, or an endurance parameter.
4. The memory sub-system of claim 1, wherein the operations comprise allocating complete portions for storing data before allocating any blocks of the incomplete portion.
5. The memory sub-system of claim 1, wherein a first set of a plurality of sets of blocks of a first portion is distributed across multiple memory dies.
6. The memory sub-system of claim 1, wherein a first set of a plurality of sets of blocks of a first portion is distributed across multiple memory planes.
7. The memory sub-system of claim 1, wherein the operations comprise storing a table, the table comprising a first entry associating a first set of a plurality of sets of blocks with a first portion address, and the table comprising a second entry associating a second set of the plurality of sets of blocks with a second portion address.
8. The memory sub-system of claim 1, wherein a total storage space corresponding to a plurality of portions including the incomplete portion corresponds to an exported capacity of the memory sub-system.
9. The memory sub-system of claim 1, wherein the one or more memory operations comprise at least one of a wear leveling operation, a data read or write operation, or a garbage collection operation.
10. The memory sub-system of claim 1, wherein the operations comprise: receiving a request to perform a memory operation; determining that each of the one or more complete portions excluding the incomplete portion has been previously allocated; and in response to determining that each of the one or more complete portions excluding the incomplete portion has been previously allocated, performing the memory operation on a set of blocks of the incomplete portion.
11. The memory sub-system of claim 10, wherein the operations comprise: allocating the incomplete portion to the memory operation.
12. The memory sub-system of claim 10, wherein the incomplete portion is a first incomplete portion, and wherein the operations comprise: determining that a first collection of blocks of a second portion of a plurality of portions is associated with an individual reliability grade that is lower than a threshold, the second portion comprising a second collection of blocks associated with

individual reliability grades that transgress the threshold; and designating the second portion as a second incomplete portion.

13. The memory sub-system of claim 12, wherein the operations comprise: determining a quantity of blocks of the first incomplete portion that are associated with a reliability grade that is lower than the threshold; selecting a group of the second collection of the blocks of the second incomplete portion that corresponds to the quantity of blocks; and associating the selected group of the second collection of blocks with the first incomplete portion.

14. The memory sub-system of claim 13, wherein the operations comprise: replacing the quantity of blocks of the first incomplete portion with the selected group of the second collection of blocks.

15. The memory sub-system of claim 14, wherein the operations comprise: performing the memory operation on the first incomplete portion in multiple passes, wherein a first address of the first incomplete portion is used in a first of the multiple passes, and wherein a second address of the second incomplete portion is used in a second of the multiple passes.

16. A computerized method comprising: designating an individual portion of a set of memory components as an incomplete portion based on a result of determining whether one or more sub-portions of the individual portion satisfy a reliability criterion; and performing one or more memory operations on an individual region of one or more complete portions of the set of memory components before the incomplete portion.

17. The computerized method of claim 16, further comprising: accessing configuration data, wherein the configuration data comprises a table that associates individual blocks of a set of memory components with respective reliability grades.

18. The computerized method of claim 17, wherein each of the respective reliability grades describes at least one of a data retention parameter, a read disturb parameter, an error rate, a leakage current, a cross temperature parameter, or an endurance parameter.

19. A non-transitory computer-readable storage medium comprising instructions that, when executed by a processing device, cause the processing device to perform operations comprising: designating an individual portion of a set of memory components as an incomplete portion based on a result of determining whether one or more sub-portions of the individual portion satisfy a reliability criterion; and performing one or more memory operations on an individual region of one or more complete portions of the set of memory components before the incomplete portion.

20. The non-transitory computer-readable storage medium of claim 19, the operations comprising accessing configuration data, wherein the configuration data comprises a table that associates individual blocks of a set of memory components with respective reliability grades.
