

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258679

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

KIM; Chansoo et al.

OPERATION METHOD OF HOST CONTROLLING COMPUTING DEVICE, AND OPERATION METHOD OF ARTIFICIAL INTELLIGENCE SYSTEM INCLUDING COMPUTING DEVICE AND HOST

Abstract

Disclosed is an operation method of a host which controls a computing device performing an artificial intelligence computation. The operation method includes receiving configuration information from the computation device, generating a plurality of lightening weight data by performing lightening on weight data based on the configuration information, generating a plurality of files by performing an aligning operation on the plurality of lightening weight data based on the configuration information, and loading the plurality of files into a memory device of the computing device. The configuration information includes channel information about a plurality of channels between the memory device and an accelerator, which are included in the computing device, and the number of the plurality of files is equal to the number of the plurality of channels.

Inventors: KIM; Chansoo (Suwon-si, KR), LEE; Hwang (Suwon-si, KR), JANG; Jae Hun (Suwon-si, KR), JEON; Younho (Suwon-si, KR), HEO; Wan (Suwon-si, KR), KWON; Sejung (Seongnam-si, KR), KIM; Byeonguk (Seongnam-si, KR), PARK; Baeseong (Seongnam-si, KR), LEE; Dongsoo (Seongnam-si, KR)

Applicant: Samsung Electronics Co., Ltd. (Suwon-si, KR)

Family ID: 96639378

Assignee: Samsung Electronics Co., Ltd. (Suwon-si, KR); NAVER CORPORATION (Seongnam-si, KR)

Appl. No.: 19/029311

Filed: January 17, 2025

Foreign Application Priority Data

KR

10-2024-0019667

Feb. 08, 2024

Publication Classification

Int. Cl.: G06F9/445 (20180101)

U.S. Cl.:

CPC G06F9/44505 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

[0001] This application claims priority under 35 U.S.C. § 119 to Korean Patent Application No. 10-2024-0019667 filed on Feb. 8, 2024, in the Korean Intellectual Property Office, the disclosures of which are incorporated by reference herein in their entireties.

BACKGROUND

[0002] Embodiments of the present disclosure described herein relate to artificial intelligence, and more particularly, relate to an operation method of a host controlling a computing device and an operation method of an artificial intelligence system including the computing device and the host.

[0003] The artificial intelligence (AI) is a field of computer science which simulates various human abilities such as learning, inference, and perception. Nowadays, artificial intelligence is widely used in various fields such as understanding natural languages, translation of natural languages, robotics, artificial vision, problem solving, learning, knowledge acquisition, and cognitive science by enabling computer systems to make inferences that would otherwise be impractical for a human to make and/or enabling the computer systems to make decisions without additional human input.

[0004] The artificial intelligence is implemented based on various algorithms. As an example, a neural network consists of a complex network in which nodes and synapses are repeatedly connected. Signal processing may be variously performed based on a relevant synapse in the process of transferring data from a current node to a next node, and this signal processing process is called a layer. That is, the neural network may include various layers which are complexly connected to each other. Because various layers included in the neural network require a large amount of computation and a large amount of data, various methods for optimizing the computation amount and the data amount are being developed.

SUMMARY

[0005] Embodiments of the present disclosure provide an operation method of a host controlling a computing device with improved performance and an operation method of an artificial intelligence system including the computing device and the host.

[0006] According to at least one embodiment, an operation method of a host configured to control a computing device performing an artificial intelligence computation includes receiving configuration information from the computation device, generating a plurality of lightening weight data by performing lightening on weight data based on the configuration information, generating a plurality of files by performing an aligning operation on the plurality of lightening weight data based on the configuration information, and loading the plurality of files into a memory device of the computing device. The configuration information includes channel information about a plurality of channels between the memory device of the computing device and an accelerator of the computing device, and the number of the plurality of files generated is equal to the number of the plurality of channels.

[0007] According to at least one embodiment, an operation method of a host configured to control a

computing device performing an artificial intelligence computation includes receiving configuration information from the computation device, generating first lightening weight data including first weight fragments and second lightening weight data including second weight fragments by performing a lightening operation on first weight data and second weight data based on the configuration information, generating a plurality of files by performing an aligning operation on the first weight fragments and the second weight fragments, and loading the plurality of files into a plurality of memories of the computing device. The configuration information includes channel information about a plurality of channels between the plurality of memories of the computing device and an accelerator of the computing device, the number of the plurality of files generated is equal to the number of the plurality of channels, the first weight fragments are distributed to the plurality of files, and the second weight fragments are distributed to the plurality of files.

[0008] According to at least one embodiment, an operation method of an artificial intelligence system which includes a computing device configured to perform an artificial intelligence computation and a host configured to communicate with the computing device, the method includes receiving, by the host, configuration information of the computation device, generating, by the host, a plurality of files by performing a lightening operation and an aligning operation on weight data based on the configuration information, loading, by the host, the plurality of files into a plurality of memories of the computing device, reading out, by the computing device, the plurality of files by performing direct memory access (DMA) operations on the plurality of memories in parallel, and performing, by the computing device, an artificial intelligence computation based on the plurality of files. The configuration information includes channel information about a plurality of channels between the plurality of memories of the computing device and an accelerator of the computing devices, and the generating, by the host, the plurality of files includes generating the plurality of files such that the number of the plurality of files is equal to the number of the plurality of channels.

Description

BRIEF DESCRIPTION OF THE FIGURES

[0009] The above and other objects and features of the present disclosure will become apparent by describing in detail embodiments thereof with reference to the accompanying drawings.

[0010] FIG. 1 is a block diagram illustrating an artificial intelligence system according to at least one embodiment of the present disclosure.

[0011] FIG. 2 is a flowchart illustrating an operation of a host of FIG. 1.

[0012] FIG. 3 is a block diagram illustrating a computing device of FIG. 1.

[0013] FIG. 4 is a diagram for describing an operation of a lightening module of a host of FIG. 1.

[0014] FIGS. 5A and 5B are diagrams for describing an operation of an aligning module of a host of FIG. 1.

[0015] FIGS. 6A and 6B are diagrams for describing an operation of an aligning module of a host of FIG. 1.

[0016] FIGS. 7A and 7B are diagrams for describing an operation of an aligning module of a host of FIG. 1.

[0017] FIG. 8 is a flowchart illustrating an operation of a computing device of FIG. 1.

[0018] FIG. 9 is a diagram for describing an operation of a lightening module of a host of FIG. 1.

[0019] FIGS. 10A and 10B are diagrams for describing an operation of an aligning module of a host of FIG. 1.

[0020] FIGS. 11A to 11C are diagrams for describing a method of reading out a plurality of files described with reference to FIGS. 10A and 10B.

[0021] FIG. 12 is a flowchart illustrating an operation of a host of FIG. 1.

[0022] FIGS. 13A and 13B are diagrams for describing an operation of an aligning module of a host of FIG. 1.

[0023] FIGS. 14A and 14B are diagrams for describing an operation of an aligning module of a host of FIG. 1.

[0024] FIG. 15 is a diagram for describing an operation of an aligning module of a host of FIG. 1.

[0025] FIG. 16 is a diagram for describing an operation of an aligning module of a host of FIG. 1.

[0026] FIG. 17 is a block diagram illustrating an artificial intelligence system according to at least one embodiment of the present disclosure.

DETAILED DESCRIPTION

[0027] Below, embodiments of the present disclosure will be described in detail and clearly to such an extent that an ordinary one in the art easily carries out the present disclosure.

[0028] In the detailed description or drawings, functional blocks which are expressed by using the terms “unit”, “element,” “module”, etc. may be implemented in the form of processing circuitry such as hardware, software, or a combination thereof, which is configured to perform a specific function. As an example, a “lightening module” may be hardware or software configured to support a data lightening operation or function disclosed in the specification, or a combination thereof. In at least some embodiments the processing circuitry more specifically may be included in (and/or enabled by), but is not limited to, a central processing unit (CPU), an arithmetic logic unit (ALU), a digital signal processor, a microcomputer, a field programmable gate array (FPGA), a System-on-Chip, (SoC), a programmable logic unit, a microprocessor, application-specific integrated circuit (ASIC), etc. and/or may include active and/or passive electrical components such as transistors, resistors, capacitors, etc., and/or electronic circuits including one or more of said components configured to perform and/or enable the function of the functional blocks.

[0029] FIG. 1 is a block diagram illustrating an artificial intelligence system according to at least one embodiment of the present disclosure. Referring to FIG. 1, an artificial intelligence system **10** may include a host **100**, an artificial intelligence computing device **200**, and a storage device **300**. In at least one embodiment, the artificial intelligence system **10** may be configured to perform various artificial intelligence-based operations such as computation, inference, or learning. The artificial intelligence system **10** may include a graphics processing unit (GPU), a neural processing unit (NPU), or separate dedicated hardware. Alternatively, the artificial intelligence system **10** may be included in an application processor (AP) and/or a mobile device. For example, artificial intelligence system **10** may be configured to perform voice recognition, text-to-speech, image recognition, image classification, and/or image processing by using a neural network, and may be included in, e.g., a tablet device, a smart TV, an augmented reality (AR) device, an Internet of things (IoT) device, a self-driving vehicle, robots, a medical device, a drone, an advanced drivers assistance system (ADAS), an image display device, a data processing server, a measuring device, etc. and/or may be mounted in one of various kinds of electronic devices.

[0030] In at least one embodiment, the host **100** may include a processor configured to perform various computations in the artificial intelligence system **10** and/or to execute various operating systems or various applications. The processor may control various components of the artificial intelligence system **10** depending on requests of the operating systems or applications. The processor may include a central processing unit (CPU) and/or an application processor (CP) which includes one or more processing cores.

[0031] The artificial intelligence computing device (for convenience of description, hereinafter referred to as a “computing device”) **200** may be configured to perform various artificial intelligence operations under control of the host **100**. For example, the computing device **200** may be configured to perform the artificial intelligence computation, inference, and/or learning based on an artificial intelligence model and/or a weight provided, e.g., from the host **100**.

[0032] In at least one embodiment, the artificial intelligence model and/or the weight may be

generated through machine learning. The machine learning may include various learning schemes such as supervised learning, unsupervised learning, semi-supervised learning, and/or reinforcement learning, but the present disclosure is not limited thereto.

[0033] In at least one embodiment, the artificial intelligence model may be generated or trained through one of various neural networks such as a deep neural network (DNN), a convolutional neural network (CNN), a recurrent neural network (RNN), a restricted boltzmann machine (RBM), a deep belief network (DBN), a bidirectional recurrent deep neural network (BRDNN), deep Q-networks, and/or a combination of two or more thereof. Alternatively, the artificial intelligence model may be a large language model (LLM) configured to perform natural language processing (NLP) and may be trained or generated by using a transformer model. The artificial intelligence model may include a plurality of neural network layers, and each of the plurality of neural network layers may be configured to perform an artificial intelligence computation based on the trained model or weight.

[0034] The storage device **300** may be configured to store various information or various data which are used in the artificial intelligence system **10**. For example, the storage device **300** may be configured to store weight data WD enabling the computing device **200** to perform the artificial intelligence computation, inference, and/or learning. In at least one embodiment, the weight data WD may be information generated by the learning of the computing device **200**. In at least one embodiment, the storage device **300** may be a high-capacity storage medium such as a solid state drive (SSD). However, the present disclosure is not limited thereto. For example, the storage device **300** may be a high-capacity storage medium such as a universal flash storage (UFS) card, an embedded UFS, or a hard disk drive (HDD) and/or may be an external storage server connected through a separate communication line.

[0035] In at least one embodiment, the weight data WD may be based on the large model, such as a large language model. In this case, the size of the weight data WD may be very large. Because the resource of the computing device **200** (e.g., the capacity of a memory device **210**) is limited. Accordingly, the host **100** may reduce the size of the weight data WD and may load the weight data WD with the reduced size into the computing device **200**, thus the entire weight data WD may not be loaded to the computing device **200** . . .

[0036] For example, the host **100** may include a lightening module **110** and an aligning module **120**. The lightening module **110** may perform lightening on the weight data WD stored in the storage device **300**. In at least one embodiment, the lightening may indicate an operation, method, and/or a technique for reducing the size and/or capacity of the weight data WD (such as pruning, quantization, compression, knowledge distillation, etc. of the weight data WD). The pruning indicates a technique for reducing the capacity of the weight data WD by removing a value less than or equal to a reference value from the weight data WD. The quantization indicates a technique for reducing the capacity of the weight data WD by converting the number of bits expressed in the weight data WD into the number of bits with relatively low precision. The compression indicates a technique for reducing the capacity of the weight data WD by compressing the weight data WD. The knowledge distillation indicates a technique for reducing the capacity of the weight data WD by training a lightening student model by using the weight data WD as a teacher model.

[0037] The above lightening techniques are provided only as some examples, and the present disclosure is not limited thereto. The lightening module **110** may be configured to generate lightening weight data LWD by lightening the weight data WD based on various lightening techniques depending on a kind of an artificial intelligence computation to be performed in the computing device **200** or a kind of an artificial intelligence model.

[0038] The aligning module **120** may be configured to generate a lightening weight data file (for convenience of description, hereinafter referred to as a “lightening weight file”) FILE_LWD by aligning the lightening weight data LWD based on configuration information CONFIG of the computing device **200**. For example, the computing device **200** may include the memory device

210 and an accelerator **220**. The accelerator **220** may be configured to communicate with the memory device **210** through a plurality of channels. The aligning module **120** may align the lightening weight data LWD based on the configuration information CONFIG such that the lightening weight file FILE_LWD is efficiently transmitted from the memory device **210** to the accelerator **220**. In at least one embodiment, the configuration information CONFIG may include information about the number of channels between the memory device **210** and the accelerator **220** included in the computing device **200** and information about a quantization type. For example, the configuration information CONFIG may include quantization information indicating, e.g., the number of weight fragments to be included in each of the plurality of lightening weight data LWD. The lightening weight data LWD may be loaded into the memory device **210**. An aligning operation of the aligning module **120** will be described in detail with reference to the following drawings.

[0039] In at least one embodiment, the accelerator **220** may include an internal buffer **221**. The accelerator **220** may store the lightening weight file FILE_LWD received from the memory device **210** in the internal buffer **221**. The accelerator **220** may be configured to generate the lightening weight data LWD based on the lightening weight file FILE_LWD stored in the internal buffer **221** and to perform an artificial intelligence computation based on the lightening weight data LWD.

[0040] As described above, the host **100** according to at least one embodiment of the present disclosure may generate the lightening weight file FILE_LWD by aligning the lightening weight data LWD based on the configuration information CONFIG of the computing device **200**. The lightening weight file FILE_LWD may be loaded into the memory device **210** of the computing device **200**, e.g., by the host **100**. In this case, the accelerator **220** of the computing device **200** may receive the lightening weight file FILE_LWD from the memory device **210** through the plurality of channels, and thus, the performance of the computing device **200** is improved.

[0041] FIG. 2 is a flowchart illustrating an operation of a host of FIG. 1. Referring to FIGS. 1 and 2, in operation S110, the host **100** may obtain or collect the configuration information CONFIG about the computing device **200**. In at least one embodiment, the host **100** may receive the configuration information CONFIG from the computing device **200** and/or the accelerator **220** of the computing device **200**. In at least one embodiment, the configuration information CONFIG may be received in an initialization operation of the computing device **200**. Alternatively, the configuration information CONFIG may be obtained from the computing device **200** in response to a request of the host **100**. In at least one embodiment, the configuration information CONFIG may include information about the number of channels between the memory device **210** and the accelerator **220** included in the computing device **200** and information about a quantization type.

[0042] In operation S120, the host **100** may read out the weight data WD from the storage device **300**.

[0043] In operation S130, the host **100** may generate the lightening weight data LWD by performing lightening on the weight data WD read from the storage device **300**. For example, the lightening module **110** of the host **100** may apply various lightening techniques to the weight data WD and/or perform various lightening operation on the weight data WD. In this case, the lightening on the weight data WD may be performed based on the configuration information CONFIG received from the computing device **200**. That is, the lightening module **110** may generate the lightening weight data LWD by performing lightening on the weight data WD such that a format requirement of the computing device **200** is satisfied. In at least one embodiment, the lightening weight data LWD may include a plurality of weight fragments. The number of weight fragments and/or the sizes of the plurality of weight fragments may be determined depending on a quantization technique (or operation) and/or a compression technique (or operation) which the computing device **200** supports.

[0044] In operation S140, the host **100** may generate the lightening weight file FILE_LWD by aligning the lightening weight data LWD based on the configuration information CONFIG. For

example, the host **100** may align the lightening weight data LWD such that the lightening weight file FILE_LWD is efficiently transmitted from the memory device **210** to the accelerator **220**. Operation **S140** will be described in detail with reference to the following drawings.

[0045] In operation **S150**, the host **100** may load the lightening weight file FILE_LWD into the memory device **210**.

[0046] FIG. **3** is a block diagram illustrating a computing device of FIG. **1**. Referring to FIGS. **1** and **3**, the computing device **200** may include the memory device **210** and the accelerator **220**.

[0047] The memory device **210** may include a plurality of memories MEM**1** to MEM**4**. The plurality of memories MEM**1** to MEM**4** may indicate memory spaces which are physically and/or logically separated and/or distinguished from each other. For example, in at least one embodiment, each of the plurality of memories MEM**1** to MEM**4** may be implemented with a separate die, a separate chip, a separate device, and/or a separate module. The plurality of memories MEM**1** to MEM**4** may be configured to communicate with the accelerator **220** through a plurality of channels CH**1** to CH**4**. In at least one embodiment, the plurality of channels CH**1** to CH**4** may indicate available communication paths which enable communication independently from each other.

[0048] The accelerator **220** may communicate with the plurality of memories MEM**1** to MEM**4** through the plurality of channels CH**1** to CH**4**, respectively. The accelerator **220** may include a plurality of memory controllers MCT**1** to MCT**4**, the internal buffer **221**, and a processing element (PE) **222**.

[0049] The plurality of memory controllers MCT**1** to MCT**4** may communicate with the plurality of memories MEM**1** to MEM**4** through the plurality of channels CH**1** to CH**4**, respectively. In at least one embodiment, the plurality of memory controllers MCT**1** to MCT**4** may read out data or files stored in the plurality of memories MEM**1** to MEM**4** through a direct memory access (DMA) operation. For example, the first memory controller MCT**1** may perform a first DMA operation on the first memory MEM**1** through the first channel CH**1**; the second memory controller MCT**2** may perform a second DMA operation on the second memory MEM**2** through the second channel CH**2**; the third memory controller MCT**3** may perform a third DMA operation on the third memory MEM**3** through the third channel CH**3**; the fourth memory controller MCT**4** may perform a fourth DMA operation on the fourth memory MEM**4** through the fourth channel CH**4**; etc.

[0050] In at least one embodiment, the plurality of memory controllers MCT**1** to MCT**4** may perform the DMA operations independently of each other or in parallel. In at least one embodiment, the direct memory access (DMA) operation of each of the plurality of memory controllers MCT**1** to MCT**4** may be initiated or set by a separate DMA controller.

[0051] The internal buffer **221** may be configured to temporarily store data or files received through the plurality of memory controllers MCT**1** to MCT**4**.

[0052] The processing element **222** may perform the artificial intelligence computation based on the data temporarily stored in the internal buffer **221**. For example, the processing element **222** may perform a multiply and accumulate (MAC) operation on weight data and activate data stored in the internal buffer **221**. In at least one embodiment, a computational result of the processing element **222** may be stored in the internal buffer **221** or the plurality of memories MEM**1** to MEM**4**.

[0053] As described above, the memory device **210** and the accelerator **220** of the computing device **200** may communicate with each other through the plurality of channels CH**1** to CH**4** to improve a transmission speed. However, because a conventional host uploads weight data to the memory device **210** regardless of a structure of the memory device **210** and the accelerator **220** included in the computing device **200**, the data transmission between the memory device **210** and the accelerator **220** is not efficiently performed. For example, the conventional host may upload lightening weight data to the memory device **210** in the form of a single file. In this case, even though the memory device **210** and the accelerator **220** communicate with each other through the plurality of channels CH**1** to CH**4**, the accelerator **220** can only access the lightening weight data from the memory device **210** through a single channel.

[0054] In contrast, according to the embodiments of the present disclosure, the aligning module **120** of the host **100** may align the lightening weight data LWD to a plurality of files based on the configuration information CONFIG of the computing device **200**. The plurality of files may be respectively loaded into the plurality of memories MEM1 to MEM4 of the memory device **210**. In this case, the accelerator **220** may access the plurality of files from the plurality of memories MEM1 to MEM4 simultaneously, independently, and/or in parallel. Accordingly, a speed at which the accelerator **220** accesses lightening weight data from the memory device **210** is improved.

[0055] FIG. **4** is a diagram for describing an operation of a lightening module of a host of FIG. **1**. Referring to FIGS. **1** and **4**, the weight data WD may include first to fourth weight data WD1, WD2, WD3, and WD4. Each of the first to fourth weight data WD1, WD2, WD3, and

[0056] WD4 may correspond to one weight value. Alternatively, each of the first to fourth weight data WD1, WD2, WD3, and WD4 may include weight values, the number of which is determined in advance. For example, each of the first to fourth weight data WD1, WD2, WD3, and WD4 may include 16 weight values. However, the present disclosure is not limited thereto.

[0057] In at least one embodiment, each of the first to fourth weight data WD1, WD2, WD3, and WD4 may include a plurality of weight values included in a corresponding artificial intelligence layer. For example, the artificial intelligence model which is used in the computing device **200** may include a plurality of layers. The plurality of layers may use different weight values. That is, weight values of the plurality of layers may be different from each other, and each of the first to fourth weight data WD1, WD2, WD3, and WD4 may include a plurality of weight values included in the corresponding artificial intelligence layer.

[0058] The above description of weight data is to describe embodiments of the present disclosure easily, and the present disclosure is not limited thereto. The number of weight data WD included in the artificial intelligence model may be variously expanded, and weight values included in one weight data may be variously changed and modified.

[0059] The lightening module **110** may perform lightening on the weight data WD based on the configuration information CONFIG and may generate the lightening weight data LWD. For example, the configuration information CONFIG may include information about a kind of quantization to be applied to the computing device **200**. Based on the information about the kind of quantization included in the configuration information CONFIG, the lightening module **110** may perform lightening on the first weight data WD1 to generate first lightening weight data LWD1, may perform lightening on the second weight data WD2 to generate second lightening weight data LWD2, may perform lightening on the third weight data WD3 to generate third lightening weight data LWD3, and may perform lightening on the fourth weight data WD4 to generate fourth lightening weight data LWD4.

[0060] For example, the first lightening weight data LWD1 may include 1a-th, 1b-th, 1c-th, and 1d-th weight fragments W1a, W1b, W1c, and W1d; the second lightening weight data LWD2 may include 2a-th, 2b-th, 2c-th, and 2d-th weight fragments W2a, W2b, W2c, and W2d; the third lightening weight data LWD3 may include 3a-th, 3b-th, 3c-th, and 3d-th weight fragments W3a, W3b, W3c, and W3d; the fourth lightening weight data LWD4 may include 4a-th, 4b-th, 4c-th, and 4d-th weight fragments W4a, W4b, W4c, and W4d; etc.

[0061] For example, the information about quantization included in the configuration information CONFIG may indicate 4-bit non-uniform quantization. In this case, the lightening module **110** may perform the 4-bit non-uniform quantization on the first to fourth weight data WD1 to WD4. A result of the 4-bit non-uniform quantization performed for each of the first to fourth weight data WD1 to WD4 may be divided into four unit data. The lightening module **110** may perform compression on the respective unit data to generate the plurality of weight fragments W1a to W4d. In this case, compression ratios of the unit data may be different from each other. That is, the plurality of weight fragments W1a to W4d may have different sizes or different lengths.

[0062] In at least one embodiment, the accelerator **220** may calculate the first weight data WD1

and/or an approximate value of the first weight data WD1 based on the first lightening weight data LWD1, may calculate the second weight data WD2 and/or an approximate value of the second weight data WD2 based on the second lightening weight data LWD2, may calculate the third weight data WD3 and/or an approximate value of the third weight data WD3 based on the third lightening weight data LWD3, may calculate the fourth weight data WD4 and/or an approximate value of the fourth weight data WD4 based on the fourth lightening weight data LWD4, etc. The accelerator 220 may perform the artificial intelligence computation based on the calculated lightening weight data LWD1 to LWD4. In this case, an appropriate value may, for example, value resulting from a conversion of the first lightening weight data LWD1 that provides a correct result in the artificial intelligence computation.

[0063] FIGS. 5A and 5B are diagrams for describing an operation of an aligning module of a host of FIG. 1. For convenience of description, the lightening weight data LWD of FIGS. 4, 5A, and 5B are provided as weight data to which the 4-bit non-uniform quantization is applied. However, this is only an example, and the present disclosure is not limited thereto.

[0064] Referring to FIGS. 1, 3, 4, 5A, and 5B, the aligning module 120 may generate a plurality of files FLa to FLd by aligning the first to fourth lightening weight data LWD1 to LWD4 based on the configuration information CONFIG. For example, the aligning module 120 may generate the a-th file FLa based on the plurality of weight fragments W1a to W1d of the first lightening weight data LWD1, may generate the b-th file FLb based on the plurality of weight fragments W2a to W2d of the second lightening weight data LWD2, may generate the c-th file FLc based on the plurality of weight fragments W3a to W3d of the third lightening weight data LWD3, and may generate the d-th file FLd based on the plurality of weight fragments W4a to W4d of the fourth lightening weight data LWD4. That is, the a-th file FLa may include the 1a-th to 1d-th weight fragments W1a to W1d, the b-th file FLb may include the 2a-th to 2d-th weight fragments W2a to W2d, the c-th file FLc may include the 3a-th to 3d-th weight fragments W3a to W3d, and the d-th file FLd may include the 4a-th to 4d-th weight fragments W4a to W4d.

[0065] Each of the plurality of files FLa to FLd may indicate a unit of data to be loaded into each of the plurality of memories MEM1 to MEM4. For example, the a-th file FLa may be loaded into the first memory MEM1, the b-th file FLb may be loaded into the second memory MEM2, the c-th file FLc may be loaded into the third memory MEM3, and the d-th file FLd may be loaded into the fourth memory MEM4.

[0066] The accelerator 220 may access the plurality of files FLa to FLd from the plurality of memories MEM1 to MEM4, may obtain the first to fourth lightening weight data LWD1 to LWD4 based on the plurality of files FLa to FLd, and may perform the artificial intelligence computation.

[0067] For example, as illustrated in FIG. 5B, the first memory controller MCT1 may perform a first DMA operation DMA1 on the first memory MEM1 to read out the a-th file FLa; the second memory controller MCT2 may perform a second DMA operation DMA2 on the second memory MEM2 to read out the b-th file FLb; the third memory controller MCT3 may perform a third DMA operation DMA3 on the third memory MEM3 to read out the c-th file FLc; the fourth memory controller MCT4 may perform a fourth DMA operation DMA4 on the fourth memory MEM4 to read out the d-th file FLd; etc.

[0068] In this case, the accelerator 220 may sequentially receive the 1a-th to 1d-th weight fragments W1a to W1d from the first memory MEM1 through the first DMA operation DMA1 and may generate the first lightening weight data LWD1 at a second point in time t2 when all the 1a-th to 1d-th weight fragments W1a to W1d are received. That is, after the second point in time t2, the accelerator 220 may start a computation which uses the first lightening weight data LWD1.

[0069] The accelerator 220 may sequentially receive the 2a-th to 2d-th weight fragments W2a to W2d from the second memory MEM2 through the second DMA operation DMA2 and may generate the second lightening weight data LWD2 at a first point in time t1 when all the 2a-th to 2d-th weight fragments W2a to W2d are received. That is, after the first point in time t1, the

accelerator **220** may start a computation which uses the second lightening weight data LWD2.

[0070] Likewise, the accelerator **220** may sequentially receive the **3a**-th to **3d**-th weight fragments **W3a** to **W3d** from the third memory **MEM3** through the third DMA operation **DMA3** and may sequentially receive the **4a**-th to **4d**-th weight fragments **W4a** to **W4d** from the fourth memory **MEM4** through the fourth DMA operation **DMA4**. At a third point in time **t3**, the accelerator **220** may generate the third lightening weight data **LWD3** and may then perform a corresponding computation; at a fourth point in time **t4**, the accelerator **220** may generate the fourth lightening weight data **LWD4** and may then perform a corresponding computation.

[0071] As described above, as the lightening weight data **LWD1**, **LWD2**, **LWD3**, and **LWD4** are distributed and stored in the plurality of memories **MEM1**, **MEM2**, **MEM3**, and **MEM4**, the accelerator **220** may receive the lightening weight data **LWD1**, **LWD2**, **LWD3**, and **LWD4** through individual DMA operations on the plurality of memories **MEM1**, **MEM2**, **MEM3**, and **MEM4**.

[0072] FIGS. **6A** and **6B** are diagrams for describing an operation of an aligning module of a host of FIG. **1**. For convenience of description, additional description associated with the components described above will be omitted to avoid redundancy. Referring to FIGS. **1**, **2**, **4**, **6A**, and **6B**, the aligning module **120** may align the lightening weight data **LWD** to generate a plurality of files **FLe**, **FLf**, **FLg**, and **FLh**.

[0073] In the example of FIG. **5A**, compression data of the same lightening weight data are included in the same file. For example, the weight fragments **W1a**, **W1b**, **W1c**, and **W1d** of the first lightening weight data **LWD1** are included in the same **a**-th file **FLa**. In contrast, in the example of FIG. **6A**, weight fragments of the same lightening weight data may be included in different files. For example, the weight fragments **W1a**, **W1b**, **W1c**, and **W1d** of the first lightening weight data **LWD1** may be respectively included in **e**-th to **h**-th files **FLe** to **FLh**; the weight fragments **W2a**, **W2b**, **W2c**, and **W2d** of the second lightening weight data **LWD2** may be respectively included in the **e**-th to **h**-th files **FLe** to **FLh**; the weight fragments **W3a**, **W3b**, **W3c**, and **W3d** of the third lightening weight data **LWD3** may be respectively included in the **e**-th to **h**-th files **FLe** to **FLh**; the weight fragments **W4a**, **W4b**, **W4c**, and **W4d** of the fourth lightening weight data **LWD4** may be respectively included in the **e**-th to **h**-th files **FLe** to **FLh**; etc.

[0074] In this case, the **e**-th file **FLe** may include the **1a**-th, **2a**-th, **3a**-th, and **4a**-th weight fragments **W1a**, **W2a**, **W3a**, and **W4a**, the **f**-th file **FLf** may include the **1b**-th, **2b**-th, **3b**-th, and **4b**-th weight fragments **W1b**, **W2b**, **W3b**, and **W4b**, the **g**-th file **FLg** may include the **1c**-th, **2c**-th, **3c**-th, and **4c**-th weight fragments **W1c**, **W2c**, **W3c**, and **W4c**, and the **h**-th file **FLh** may include the **1d**-th, **2d**-th, **3d**-th, and **4d**-th weight fragments **W1d**, **W2d**, **W3d**, and **W4d**. The **e**-th to **h**-th files **FLe** to **FLh** may be respectively stored in the first to fourth memories **MEM1** to **MEM4**.

[0075] In at least one embodiment, as compression data of the same lightening weight data are respectively stored in different files, a point in time when the artificial intelligence computation of the accelerator **220** starts may be advanced. For example, as illustrated in FIG. **6B**, the accelerator **220** may read out the **e**-th to **h**-th files **FLe** to **FLh** from the first to fourth memories **MEM1** to **MEM4** in parallel.

[0076] For example, the first memory controller **MCT1** may sequentially read out the weight fragments **W1a**, **W2a**, **W3a**, and **W4a** of the **e**-th file **FLe** from the first memory **MEM1** through the first DMA operation **DMA1**. The second memory controller **MCT2** may sequentially read out the weight fragments **W1b**, **W2b**, **W3b**, and **W4b** of the **h**-th file **FLh** from the second memory **MEM2** through the second DMA operation **DMA2**. The third memory controller **MCT3** may sequentially read out the weight fragments **W1c**, **W2c**, **W3c**, and **W4c** of the **f**-th file **FLf** from the third memory **MEM3** through the third DMA operation **DMA3**. The fourth memory controller **MCT4** may sequentially read out the weight fragments **W1d**, **W2d**, **W3d**, and **W4d** of the **g**-th file **FLg** from the fourth memory **MEM4** through the fourth DMA operation **DMA4**.

[0077] In this case, at a first point in time **t1** when all the weight fragments **W1a**, **W1b**, **W1c**, and **W1d** corresponding to the first lightening weight data **LWD1** are received, the accelerator **220** may

obtain the first lightening weight data LWD1 based on the weight fragments W1a, W1b, W1c, and W1d and may start the artificial intelligence computation on the first lightening weight data LWD1. Next, at a second point in time t2 when all the weight fragments W2a, W2b, W2c, and W2d corresponding to the second lightening weight data LWD2 are received, the accelerator 220 may obtain the second lightening weight data LWD2 based on the weight fragments W2a, W2b, W2c, and W2d and may start the artificial intelligence computation on the second lightening weight data LWD2. Then, at a third point in time t3 when all the weight fragments W3a, W3b, W3c, and W3d corresponding to the third lightening weight data LWD3 are received, the accelerator 220 may obtain the third lightening weight data LWD3 based on the weight fragments W3a, W3b, W3c, and W3d and may start the artificial intelligence computation on the third lightening weight data LWD3. Finally, at a fourth point in time t4 when all the weight fragments W4a, W4b, W4c, and W4d corresponding to the fourth lightening weight data LWD4 are received, the accelerator 220 may obtain the fourth lightening weight data LWD4 based on the weight fragments W4a, W4b, W4c, and W4d and may start the artificial intelligence computation on the fourth lightening weight data LWD4.

[0078] As described above, the aligning module 120 may perform the aligning operation based on the configuration information CONFIG of the computing device 200 such that weight fragments of the lightening weight data LWD are distributed and included in a plurality of files. In this case, the accelerator 220 of the computing device 200 may access a plurality of files from the plurality of memories MEM1 to MEM4 in parallel and may obtain lightening weight data based on the plurality of files accessed in parallel.

[0079] FIGS. 7A and 7B are diagrams for describing an operation of an aligning module of a host of FIG. 1. For convenience of description, additional description associated with the components described above will be omitted to avoid redundancy. Referring to FIGS. 1, 2, 4, 7A, and 7B, the aligning module 120 may align the lightening weight data LWD to generate a plurality of files FL1, FL2, FL3, and FL4.

[0080] In the embodiment of FIG. 7A, weight fragments of the same lightening weight data are included in different files. For example, the weight fragments W1a, W1b, W1c, and W1d of the first lightening weight data LWD1 may be respectively included in the first to fourth files FL1 to FL4; the weight fragments W2a, W2b, W2c, and W2d of the second lightening weight data LWD2 may be respectively included in the first to fourth files FL1 to FL4; the weight fragments W3a, W3b, W3c, and W3d of the third lightening weight data LWD3 may be respectively included in the first to fourth files FL1 to FL4. The weight fragments W4a, W4b, W4c, and W4d of the fourth lightening weight data LWD4 may be respectively included in the first to fourth files FL1 to FL4; etc.

[0081] In this case, the first file FL1 may include the weight fragments W1a, W2a, W3a, and W4a, the second file FL2 may include the weight fragments W1b, W2b, W3b, and W4b, the third file FL3 may include the weight fragments W1c, W2c, W3c, and W4c, and the fourth file FL4 may include the weight fragments W1d, W2d, W3d, and W4d.

[0082] In the embodiment of FIG. 7A, the aligning module 120 may add padding data such that weight fragments included in the first to fourth files FL1 to FL4 have the same length.

[0083] For example, in the embodiment of FIGS. 6A and 6B, because weight fragments included in the plurality of files FL_e to FL_h have different lengths, individual DMA setting is required for each of the first to fourth memory controllers MCT1 to MCT4. In at least one embodiment, the DMA setting may be performed by a separate DMA controller included in the computing device 200.

[0084] For example, for the first to fourth memory controllers MCT1 to MCT4 to perform the DMA operations, the setting of DMA start addresses, offsets, etc. is required. In this case, in the example of FIGS. 6A and 6B, when the first memory controller MCT1 accesses the 2a-th weight fragment W2a from the first memory MEM1, a first start address, a first offset, etc. for the 2a-th weight fragment W2a are set in the first memory controller MCT1. Also, when the second memory

controller MCT2 accesses the $2b$ -th weight fragment $W2b$ from the second memory MEM2, a second start address, a second offset, etc. for the $2b$ -th weight fragment $W2b$ are set in the second memory controller MCT2. In this case, as illustrated in FIGS. 6A and 6B, the size of the weight fragment included in the e -th file FL_e is different from the size of the weight fragment included in the f -th file FL_f , the first and second start addresses may be different from each other, and the first and second offsets may be different from each other. Also, a point in time when the first memory controller MCT1 accesses the $2a$ -th weight fragment $W2a$ may be different from a point in time when the second memory controller MCT2 accesses the $2b$ -th weight fragment $W2b$. In this case, the DMA setting should be performed for each of the first memory controller MCT1 and the second memory controller MCT2, which causes the increase in complexity of the computing device 200.

[0085] In contrast, according to the embodiment of FIGS. 7A and 7B, the aligning module 120 may add padding data such that weight fragments included in the first to fourth files $FL1$ to $FL4$ have the same length. For example, the aligning module 120 may add padding data to the weight fragments $W1a$ to $W4d$ of the first to fourth lightening weight data $LWD1$ to $LWD4$ such that the weight fragments $W1a$ to $W4d$ of the first to fourth lightening weight data $LWD1$ to $LWD4$ are stored at the same address in the plurality of memories MEM1 to MEM4.

[0086] In detail, as illustrated in FIG. 7A, the first file $FL1$ may include the $1a$ -th and $2a$ -th weight fragments $W1a$ and $W2a$, the second file $FL2$ may include the $1b$ -th and $2b$ -th weight fragments $W1b$ and $W2b$, the third file $FL3$ may include the $1c$ -th and $2c$ -th weight fragments $W1c$ and $W2c$, and the fourth file $FL4$ may include the $1d$ -th and $2d$ -th weight fragments $W1d$ and $W2d$.

[0087] In this case, the aligning module 120 may add padding data to the remaining weight fragments $W1b$, $W1c$, and $W1d$ based on the first $1a$ -th weight fragment $W1a$ having the largest size or the longest length from among the $1a$ -th, $1b$ -th, $1c$ -th, and $1d$ -th weight fragments $W1a$, $W1b$, $W1c$, and $W1d$.

[0088] Likewise, the first file $FL1$ may further include the $3a$ -th and $4a$ -th weight fragments $W3a$ and $W4a$, the second file $FL2$ may further include the $3b$ -th and $4b$ -th weight fragments $W3b$ and $W4b$, the third file $FL3$ may further include the $3c$ -th and $4c$ -th weight fragments $W3c$ and $W4c$, and the fourth file $FL4$ may further include the $3d$ -th and $4d$ -th weight fragments $W3d$ and $W4d$. The aligning module 120 may add padding data to the remaining weight fragments $W2b$ and $W2c$ based on the $2a$ -th or $2d$ -th weight fragments $W2a$ or $W2d$ having the largest size or the longest length from among the $2a$ -th, $2b$ -th, $2c$ -th, and $2d$ -th weight fragments $W2a$, $W2b$, $W2c$, and $W2d$, may add padding data to the remaining weight fragments $W3b$, $W3c$, and $W3d$ based on the $3a$ -th weight fragment $W3a$ having the largest size or the longest length from among the $3a$ -th, $3b$ -th, $3c$ -th, and $3d$ -th weight fragments $W3a$, $W3b$, $W3c$, and $W3d$, and may add padding data to the remaining weight fragments $W4a$, $W4b$, and $W4d$ based on the $4c$ -th weight fragment $W4c$ having the largest size or the longest length from among the $4a$ -th, $4b$ -th, $4c$ -th, and $4d$ -th weight fragments $W4a$, $W4b$, $W4c$, and $W4d$.

[0089] In this case, the first to fourth files $FL1$ to $FL4$ may be aligned as illustrated in FIG. 7A, and the first to fourth files $FL1$ to $FL4$ may be respectively loaded into the first to fourth memories MEM1 to MEM4.

[0090] The first to fourth files $FL1$ to $FL4$ may be generated by aligning the lightening weight data LWD as illustrated in FIG. 7A. In this case, the DMA setting of the accelerator 220 associated with the first to fourth memory controllers MCT1 to MCT4 may be simplified.

[0091] For example, the weight fragments $W1a$, $W1b$, $W1c$, and $W1d$ corresponding to the first lightening weight data $LWD1$ may be respectively read out from the first to fourth memories MEM1 to MEM4 by the first to fourth memory controllers MCT1 to MCT4. In this case, as padding data are added to some of the weight fragments $W1a$, $W1b$, $W1c$, and $W1d$, the weight fragments $W1a$, $W1b$, $W1c$, and $W1d$ may have the same offsets. Accordingly, the same first start addresses and the same first offsets may be set in the first to fourth memory controllers MCT1 to

MCT4, and thus, the weight fragments W1a, W1b, W1c, and W1d may be read out from the first to fourth memories MEM1 to MEM4 during a time interval from t0 to t1.

[0092] The weight fragments W2a, W2b, W2c, and W2d corresponding to the second lightening weight data LWD2 may be respectively read out from the first to fourth memories MEM1 to MEM4 by the first to fourth memory controllers MCT1 to MCT4. In this case, as padding data are added to some of the weight fragments W2a, W2b, W2c, and W2d, the weight fragments W2a, W2b, W2c, and W2d may have the same offsets. Also, because the offsets corresponding to the previous weight fragments W1a, W1b, W1c, and W1d are identical to each other, start addresses of the weight fragments W2a, W2b, W2c, and W2d may be identical to each other. Accordingly, the same second start addresses and the same second offsets may be set in the first to fourth memory controllers MCT1 to MCT4, and thus, the weight fragments W2a, W2b, W2c, and W2d may be read out from the first to fourth memories MEM1 to MEM4 during a time interval from t1 to t2.

[0093] The weight fragments W3a, W3b, W3c, and W3d corresponding to the third lightening weight data LWD3 may be respectively read out from the first to fourth memories MEM1 to MEM4 by the first to fourth memory controllers MCT1 to MCT4. In this case, as padding data are added to some of the weight fragments W3a, W3b, W3c, and W3d, the weight fragments W3a, W3b, W3c, and W3d may have the same offsets. Also, because the offsets corresponding to the previous weight fragments W2a, W2b, W2c, and W2d are identical to each other, start addresses of the weight fragments W3a, W3b, W3c, and W3d may be identical to each other. Accordingly, the same third start addresses and the same third offsets may be set in the first to fourth memory controllers MCT1 to MCT4, and thus, the weight fragments W3a, W3b, W3c, and W3d may be read out from the first to fourth memories MEM1 to MEM4 during a time interval from t2 to t3.

[0094] Likewise, the weight fragments W4a, W4b, W4c, and W4d corresponding to the fourth lightening weight data LWD4 may be respectively read out from the first to fourth memories MEM1 to MEM4 by the first to fourth memory controllers MCT1 to MCT4. In this case, as padding data are added to some of the weight fragments W4a, W4b, W4c, and W4d, the weight fragments W4a, W4b, W4c, and W4d may have the same offsets. Also, because the offsets corresponding to the previous weight fragments W3a, W3b, W3c, and W3d are identical to each other, start addresses of the weight fragments W4a, W4b, W4c, and W4d may be identical to each other. Accordingly, the same fourth start addresses and the same fourth offsets may be set in the first to fourth memory controllers MCT1 to MCT4, and thus, the weight fragments W4a, W4b, W4c, and W4d may be read out from the first to fourth memories MEM1 to MEM4 during a time interval from t3 to t4.

[0095] As described above, according to at least one embodiment of the present disclosure, the aligning module 120 of the host 100 may generate the lightening weight file FILE_LWD by aligning the lightening weight data LWD based on the configuration information CONFIG of the computing device 200. In this case, in the same DMA period, the aligning module 120 may add padding data to weight fragments such that data to be respectively accessed from the plurality of memories MEM1 to MEM4 have the same lengths. Accordingly, the first to fourth memory controllers MCT1 to MCT4 of the accelerator 220 may normally read out the weight fragments W1a to W1d from the first to fourth memories MEM1 to MEM4 by using the same start addresses and the same offsets.

[0096] In at least one embodiment, the accelerator 220 may further include a separate device and/or structure for removing the padding data read out from the first to fourth memories MEM1 to MEM4.

[0097] Below, for convenience of description, the term “DMA period” is used. The DMA period indicates a period of a DMA operation which is performed by a memory controller through one setting for a start address and an offset. For example, in the embodiment of FIG. 7B, an A-th DMA period may indicate a period from t0 to t1. That is, in the A-th DMA period, the first memory controller MCT1 may read out the 1a-th weight fragment W1a from the first memory MEM1, the

second memory controller MCT2 may read out the **1b**-th weight fragment **W1b** from the second memory MEM2, the third memory controller MCT3 may read out the **1c**-th weight fragment **Wc** from the third memory MEM3, and the fourth memory controller MCT4 may read out the **1d**-th weight fragment **W1d** from the fourth memory MEM4. A **B**-th DMA period may be a time period from a time interval from **t1** to **t2**. That is, in the **B**-th DMA period, the first memory controller MCT1 may read out the **2a**-th weight fragment **W2a** from the first memory MEM1, the second memory controller MCT2 may read out the **2b**-th weight fragment **W2b** from the second memory MEM2, the third memory controller MCT3 may read out the **2c**-th weight fragment **W2c** from the third memory MEM3, and the fourth memory controller MCT4 may read out the **2d**-th weight fragment **W2d** from the fourth memory MEM4.

[0098] As described above, the aligning module **120** may add padding data to weight fragments such that data to be respectively read out from the first to fourth memories MEM1 to MEM4 in the same DMA period have the same sizes. Accordingly, the performance of the computing device **200** is improved.

[0099] FIG. **8** is a flowchart illustrating an operation of a computing device of FIG. **1**. Referring to FIGS. **1** and **8**, in operation **S210**, the computing device **200** may perform the DMA operation to read out weight fragments. For example, based on the method described with reference to FIGS. **1** to **7B**, the host **100** may generate a plurality of files and may load the plurality of files thus generated into the plurality of memories MEM1 to MEM4 of the memory device **210**. The accelerator **220** of the computing device **200** may perform the DMA operation on each of the plurality of memories MEM1 to MEM4 to read out weight fragments from the plurality of memories MEM1 to MEM4.

[0100] In operation **S220**, the computing device **200** may combine the weight fragments to generate weight data, lightening weight data, and/or approximate weight data. For example, as described with reference to FIGS. **1** to **7B**, the first to fourth memory controllers MCT1 to MCT4 of the accelerator **220** may perform the DMA operations on the first to fourth memories MEM1 to MEM4 to read out weight fragments (e.g., **W1a** to **W4d**). The weight fragments (e.g., **W1a** to **W4d**) may be temporarily stored in the internal buffer **221** of the accelerator **220**. The accelerator **220** may generate the lightening weight data **LWD1** to **LWD4** based on the temporarily stored weight fragments (e.g., **W1a** to **W4d**) in the internal buffer **221**.

[0101] For example, the accelerator **220** may combine the **1a**-th, **1b**-th, **1c**-th, and **1d**-th weight fragments **W1a**, **W1b**, **W1c**, and **W1d** to generate the first lightening weight data **LWD1**, may combine the **2a**-th, **2b**-th, **2c**-th, and **2d**-th weight fragments **W2a**, **W2b**, **W2c**, and **W2d** to generate the second lightening weight data **LWD2**, may combine the **3a**-th, **3b**-th, **3c**-th, and **3d**-th weight fragments **W3a**, **W3b**, **W3c**, and **W3d** to generate the third lightening weight data **LWD3**, may combine the **4a**-th, **4b**-th, **4c**-th, and **4d**-th weight fragments **W4a**, **W4b**, **W4c**, and **W4d** to generate the fourth lightening weight data **LWD4**.

[0102] In operation **S230**, the computing device **200** may perform the artificial intelligence computation by using the weight data. For example, the computing device **200** may perform the MAC operation by using the weight data. In at least one embodiment, a computational result of the computing device **200** may be again stored in the memory device **210**.

[0103] FIG. **9** is a diagram for describing an operation of a lightening module of a host of FIG. **1**. Referring to FIGS. **1** and **9**, the weight data **WD** may include the first to fourth weight data **WD1**, **WD2**, **WD3**, and **WD4**. The weight data **WD** is described with reference to FIG. **4**. In other words, the weight data **WD** of FIG. **9** may be the same as and/or substantially similar to the weight data **WD** of FIG. **4**, thus, additional description will be omitted to avoid redundancy.

[0104] The lightening module **110** may perform lightening on the weight data **WD** based on the configuration information **CONFIG** and may generate lightening weight data **LWD-1**. For example, the configuration information **CONFIG** may include information about a kind of quantization to be applied to the computing device **200**. Based on the information about the kind of quantization

included in the configuration information CONFIG, the lightening module **110** may perform lightening on the first weight data WD1 to generate first lightening weight data LWD1-1, may perform lightening on the second weight data WD2 to generate second lightening weight data LWD2-1, may perform lightening on the third weight data WD3 to generate third lightening weight data LWD3-1, and may perform lightening on the fourth weight data WD4 to generate fourth lightening weight data LWD4-1. For example, the information about quantization included in the configuration information CONFIG may indicate 3-bit non-uniform quantization. In this case, the lightening module **110** may perform the 3-bit non-uniform quantization on the first to fourth weight data WD1 to WD4. A result of the 3-bit non-uniform quantization performed for each of the first to fourth weight data WD1 to WD4 may be divided into three data chunks. The lightening module **110** may perform compression on the respective data chunks to generate the plurality of weight fragments W1a to W4c. In this case, compression ratios of the data chunks may be different from each other. That is, the plurality of weight fragments W1a to W4c may have different sizes and/or different lengths.

[0105] In at least one embodiment, the first lightening weight data LWD1-1 may include the 1a-th, 1b-th, and 1c-th weight fragments W1a, W1b, and W1c; the second lightening weight data LWD2-1 may include the 2a-th, 2b-th, and 2c-th weight fragments W2a, W2b, and W2c; the third lightening weight data LWD3-1 may include the 3a-th, 3b-th, and 3c-th weight fragments W3a, W3b, and W3c; and the fourth lightening weight data LWD4-1 may include the 4a-th, 4b-th, and 4c-th weight fragments W4a, W4b, and W4c.

[0106] FIGS. **10A** and **10B** are diagrams for describing an operation of an aligning module of a host of FIG. **1**. For convenience of description, the lightening weight data LWD-1 of FIGS. **9**, **10A**, and **10B** are provided as the weight data to which the 3-bit non-uniform quantization is applied. However, this is only an example, and the present disclosure is not limited thereto.

[0107] Referring to FIGS. **1**, **3**, **9**, **10A**, and **10B**, the aligning module **120** may generate a plurality of files FL1-1 to FL4-1 by aligning the first to fourth lightening weight data LWD1-1 to LWD4-1 based on the configuration information CONFIG.

[0108] For example, the configuration information CONFIG may include information about four channels CH1 to CH4 between the memory device **210** and the accelerator **220** of the computing device **200**. The aligning module **120** may generate the four files FL1-1 to FL4-1 based on the configuration information CONFIG (i.e., information of the four channels CH1 to CH4). In this case, the aligning module **120** may distribute the weight fragments W1a to W4c of the first to fourth lightening weight data LWD1-1 to LWD4-1 into the four files FL1-1 to FL4-1. The four files FL1-1 to FL4-1 may be respectively loaded into the first to fourth memories MEM1 to MEM4.

[0109] In at least one embodiment, the first lightening weight data LWD1-1 may include the three weight fragments W1a to W1c, the second lightening weight data LWD2-1 may include the three weight fragments W2a to W2c, the third lightening weight data LWD3-1 may include the three weight fragments W3a to W3c, and the fourth lightening weight data LWD4-1 may include the three weight fragments W4a to W4c. That is, the number of weight fragments may be 12; in this case, each of the first to fourth files FL1-1 to FL4-1 may include three weight fragments. For example, the first file FL1-1 may include the 1a-th, 2b-th, and 3c-th weight fragments W1a, W2b, and W3c, the second file FL2-1 may include the 1b-th, 2c-th, and 4a-th weight fragments W1b, W2c, and W4a, the third file FL3-1 may include the 1c-th, 3a-th, and 4b-th weight fragments W1c, W3a, and W4b, and the fourth file FL4-1 may include 2a-th, 3b-th, and 4c-th weight fragments W2a, W3b, and W4c.

[0110] In at least one embodiment, padding data may be added to weight fragments included in each of the first to fourth files FL1-1 to FL4-1, based on the method described with reference to FIGS. **7A** and **7B**. For example, in the weight fragments W1a to W4c included in the first to fourth files FL1-1 to FL4-1, based on the 1a-th weight fragment W1a having the longest length or the largest size from among the 1a-th, 1b-th, 1c-th, and 2a-th weight fragments W1a, W1b, W1c, and

W2a to be accessed during the A-th DMA period, padding data may be added to the remaining weight fragments W1b, W1c, and W2a. Accordingly, the 1a-th, 1b-th, 1c-th, and 2a-th weight fragments W1a, W1b, W1c, and W2a to be accessed during the A-th DMA period may have the same sizes.

[0111] The first to fourth memory controllers MCT1 to MCT4 of the accelerator 220 may perform the DMA operations on the first to fourth memories MEM1 to MEM4 to access the weight fragments W1a to W4c. For example, during the A-th DMA period from t0 to t1, the accelerator 220 may read out the 1a-th, 1b-th, 1c-th, and 2a-th weight fragments W1a, W1b, W1c, and W2a from the first to fourth memories MEM1 to MEM4. During the B-th DMA period from t1 to t2, the accelerator 220 may read out the 2b-th, 2c-th, 3a-th, and 3b-th weight fragments W2b, W2c, W3a, and W3b from the first to fourth memories MEM1 to MEM4. During the C-th DMA period from t2 to t3, the accelerator 220 may read out the 3c-th, 4a-th, 4b-th, and 4c-th weight fragments W3c, W4a, W4b, and W4c from the first to fourth memories MEM1 to MEM4.

[0112] FIGS. 11A to 11C are diagrams for describing a method of reading out a plurality of files described with reference to FIGS. 10A and 10B. For convenience of description, it is assumed that the first to fourth files FL1-1 to FL4-1 described with reference to FIGS. 10A and 10B are respectively loaded into the plurality of memories MEM1 to MEM4. That is, the first memory MEM1 may include the first file FL1-1 including the 1a-th, 2b-th, and 3c-th weight fragments W1a, W2b, and W3c, the second memory MEM2 may include the second file FL2-1 including the 1b-th, 2c-th, and 4a-th weight fragments W1b, W2c, and W4a, the third memory MEM3 may include the third file FL3-1 including the 1c-th, 3a-th, and 4b-th weight fragments W1c, W3a, and W4b, and the fourth memory MEM4 may include the fourth file FL4-1 including 2a-th, 3b-th, and 4c-th weight fragments W2a, W3b, and W4c.

[0113] First, referring to FIGS. 10A, 10B, and 11A, during the A-th DMA period from t0 to t1, the first memory controller MCT1 may read out the 1a-th weight fragment W1a from the first memory MEM1, the second memory controller MCT2 may read out the 1b-th weight fragment W1b from the second memory MEM2, the third memory controller MCT3 may read out the 1c-th weight fragment W1c from the third memory MEM3, and the fourth memory controller MCT4 may read out the 2a-th weight fragment W2a from the fourth memory MEM4. The accelerator 220 may generate the first lightning weight data LWD1-1 based on the 1a-th, 1b-th, and 1c-th weight fragments W1a, W1b, and W1c. The weight fragment W2a may be maintained in the internal buffer 221 of the accelerator 220.

[0114] Next, referring to FIGS. 10A, 10B, and 11B, during the B-th DMA period from t1 to t2, the first memory controller MCT1 may read out the 2b-th weight fragment W2b, the second memory controller MCT2 may read out the 2c-th weight fragment W2c, the third memory controller MCT3 may read out the 3a-th weight fragment W3a, and the fourth memory controller MCT4 may read out the 3b-th weight fragment W3b. The accelerator 220 may generate the second lightning weight data LWD2-1 based on the 2a-th weight fragment W2a read out by the fourth memory controller MCT4 in the A-th DMA period and the 2b-th and 2c-th weight fragments W2b and W2c respectively read out by the first and second memory controllers MCT1 and MCT2 in the B-th DMA period. The 3a-th and 3b-th weight fragments W3a and W3b may be maintained in the internal buffer 221 of the accelerator 220.

[0115] Then, referring to FIGS. 10A, 10B, and 11C, during the C-th DMA period from t2 to t3, the first memory controller MCT1 may read out the 3c-th weight fragment W3c, the second memory controller MCT2 may read out the 4a-th weight fragment W4a, the third memory controller MCT3 may read out the 4b-th weight fragment W4b, and the fourth memory controller MCT4 may read out the 4c-th weight fragment W4c. The accelerator 220 may generate the third lightning weight data LWD3-1 based on the 3a-th and 3b-th weight fragment W3a and W3b read out by the third and fourth memory controllers MCT3 and MCT4 in the B-th DMA period and the 3c-th weight fragment W3c read out by the first memory controller MCT1 in the C-th DMA period. The

accelerator **220** may generate the fourth lightening weight data **LWD4-1** based on the **4a**-th to **4c**-th weight fragments **W4a** to **W4c** respectively read out by the second to fourth memory controllers **MCT2** to **MCT4** in the **C**-th DMA period.

[0116] As described above, according to at least one embodiment of the present disclosure, the host **100** may generate the lightening weight file **FILE_LWD** by performing lightening and aligning on the weight data **WD** based on the configuration information **CONFIG** of the computing device **200**. The lightening weight file **FILE_LWD** may be loaded into the plurality of memories **MEM1** to **MEM4** of the computing device **200**. In this case, the accelerator **220** of the computing device **200** may read out the lightening weight file **FILE_LWD** by the DMA operations on the plurality of memories **MEM1** to **MEM4** in parallel or individually and may obtain lightening weight data by using weight fragments included in the lightening weight file **FILE_LWD**. Accordingly, because a speed at which the accelerator **220** of the computing device **200** accesses lightening weight data is improved, the performance of the computing device **200** is improved.

[0117] FIG. **12** is a flowchart illustrating an operation of a host of FIG. **1**. Referring to FIGS. **1** and **12**, in operation **S310**, the host **100** may obtain the configuration information **CONFIG** from the computing device **200**. In operation **S320**, the host **100** may read out the weight data **WD** from the storage device **300**. In operation **S330**, the host **100** may perform lightening on the weight data **WD** to generate the lightening weight data **LWD**. Operation **S310** to operation **S330** are similar to operation **S110** to operation **S130** described with reference to FIG. **2**, and thus, additional description will be omitted to avoid redundancy.

[0118] In operation **S340**, the host **100** may generate the lightening weight file **FILE_LWD** by aligning the lightening weight data **LWD** based on the configuration information **CONFIG** and the size of a weight fragment. For example, based on the configuration information **CONFIG** and the size of the weight fragment, the aligning module **120** of the host **100** may generate the lightening weight file **FILE_LWD** such that the size of the lightening weight file **FILE_LWD** is minimized. Alternatively, based on the configuration information **CONFIG** and the size of the weight fragment, the aligning module **120** of the host **100** may generate the lightening weight file **FILE_LWD** such that of the lightening weight file **FILE_LWD** is efficiently transmitted from the memory device **210** to the accelerator **220**. Operation **S340**, that is, the aligning operation will be described in detail with reference to FIGS. **13A** to **16**.

[0119] In operation **S350**, the host **100** may load the lightening weight file **FILE_LWD** into the memory device **210**.

[0120] FIGS. **13A** and **13B** are diagrams for describing an operation of an aligning module of a host of FIG. **1**. In at least one embodiment, operation **S340** of FIG. **12** will be described with reference to FIGS. **13A** and **13B**. For convenience of description, repeat description associated with the components already described above will be omitted to avoid redundancy. Referring to FIGS. **1**, **12**, **13A**, and **13B**, the aligning module **120** of the host **100** may perform the aligning operation on the lightening weight data **LWD**, based on the configuration information **CONFIG**.

[0121] For example, the lightening weight data **LWD** may be a result of performing the 4-bit non-uniform quantization and the compression on the weight data **WD**. That is, the first lightening weight data **LWD1** may include the **1a**-th to **1d**-th weight fragments **W1a**, **W1b**, **W1c**, and **W1d**, the second lightening weight data **LWD2** may include the **2a**-th to **2d**-th weight fragments **W2a** to **W2d**, the third lightening weight data **LWD3** may include the **3a**-th to **3d**-th weight fragments **W3a** to **W3d**, and the fourth lightening weight data **LWD4** may include the **4a**-th to **4d**-th weight fragments **W4a** to **W4d**.

[0122] The aligning module **120** may place the **1a**-th to **1d**-th weight fragments **W1a**, **W1b**, **W1c**, and **W1d** of the first lightening weight data **LWD1** in first to fourth files **FL1-2** to **FL4-2**, respectively. Afterwards, the aligning module **120** may place the **2a**-th to **2d**-th weight fragments **W2a** to **W2d** of the second lightening weight data **LWD2** in the first to fourth files **FL1-2** to **FL4-2**, respectively, based on sizes of the **1a**-th to **1d**-th weight fragments **W1a**, **W1b**, **W1c**, and **W1d** of

the first lightening weight data LWD1 and sizes of the 2a-th to 2d-th weight fragments W2a to W2d of the second lightening weight data LWD2.

[0123] For example, the 1a-th weight fragment W1a among the 1a-th to 1d-th weight fragments W1a, W1b, W1c, and W1d may have the longest length or the largest size. In this case, the 2b-th weight fragment W2b having the shortest length or the smallest size from among the 2a-th to 2d-th weight fragments W2a to W2d may be added to the first file FL1-2 such that the 2b-th weight fragment W2b is placed after the 1a-th weight fragment W1a.

[0124] The 1c-th weight fragment W1c among the 1a-th to 1d-th weight fragments W1a, W1b, W1c, and W1d may have the second longest length or the second largest size. In this case, the 2c-th weight fragment W2c having the second shortest length or the second smallest size from among the 2a-th to 2d-th weight fragments W2a to W2d may be added to the second file FL3-2 such that the 2c-th weight fragment W2c is placed after the 1c-th weight fragment W1c.

[0125] The 1b-th weight fragment W1b among the 1a-th to 1d-th weight fragments W1a, W1b, W1c, and W1d may have the third longest length or the third largest size. In this case, the 2d-th weight fragment W2d having the third shortest length or the third smallest size from among the 2a-th to 2d-th weight fragments W2a to W2d may be added to the second file FL2-2 such that the 2d-th weight fragment W2d is placed after the 1b-th weight fragment W1b.

[0126] The 1d-th weight fragment W1d among the 1a-th to 1d-th weight fragments W1a, W1b, W1c, and W1d may have the shortest length or the smallest size. In this case, the 2a-th weight fragment W2a having the longest length or the largest size from among the 2a-th to 2d-th weight fragments W2a to W2d may be added to the fourth file FL4-2 such that the 2a-th weight fragment W2a is placed after the 1d-th weight fragment W1d.

[0127] Likewise, the aligning module 120 may place the 3a-th to 3d-th weight fragments W3a to W3d of the third lightening weight data LWD3 and the 4a-th to 4d-th weight fragments W4a to W4d of the fourth lightening weight data LWD4 in the first to fourth files FL1-2 to FL4-2, respectively, based on sizes of the 3a-th to 3d-th weight fragments W3a to W3d of the third lightening weight data LWD3 and sizes of the 4a-th to 4d-th weight fragments W4a to W4d of the fourth lightening weight data LWD4.

[0128] According to the above description, the first file FL1-2 may include the 1a-th, 2b-th, 3c-th, and 4a-th weight fragments W1a, W2b, W3c, and W4a, the second file FL2-2 may include the 1b-th, 2d-th, 3b-th, and 4d-th weight fragments W1b, W2d, W3b, and W4d, the third file FL3-2 may include the 1c-th, 2c-th, 3d-th, and 4c-th weight fragments W1c, W2c, W3d, and W4c, and the fourth file FL4-2 may include the 1d-th, 2a-th, 3a-th, and 4b-th weight fragments W1d, W2a, W3a, and W4b.

[0129] The first to fourth files FL1-2 to FL4-2 may be respectively loaded into the first to fourth memories MEM1 to MEM4. The accelerator 220 may read out the first to fourth files FL1-2 to FL4-2 or weight fragments from the first to fourth memories MEM1 to MEM4.

[0130] For example, as illustrated in FIG. 13B, the first memory controller MCT1 may sequentially receive the 1a-th, 2b-th, 3c-th, and 4a-th weight fragments W1a, W2b, W3c, and W4a from the first memory MEM1 through the first DMA operation, and the second memory controller MCT2 may sequentially receive the 1b-th, 2d-th, 3b-th, and 4d-th weight fragments W1b, W2d, W3b, and W4d from the second memory MEM2 through the second DMA operation, the third memory controller MCT3 may sequentially receive the 1c-th, 2c-th, 3d-th, and 4c-th weight fragments W1c, W2c, W3d, and W4c from the third memory MEM3 through the third DMA operation, and the fourth memory controller MCT4 may sequentially receive the 1d-th, 2a-th, 3a-th, and 4b-th weight fragments W1d, W2a, W3a, and W4b from the fourth memory MEM4 through the fourth DMA operation.

[0131] At the first point in time t1, the accelerator 220 may generate the first lightening weight data LWD1; at the second point in time t2, the accelerator 220 may generate the second lightening weight data LWD2; at the third point in time t3, the accelerator 220 may generate the third

lightening weight data LWD3; and, at the fourth point in time **t4**, the accelerator **220** may generate the fourth lightening weight data LWD4.

[0132] In at least one embodiment, as illustrated in FIGS. **13A** and **13B**, the plurality of files FL1-2 to FL4-2 may be generated based on sizes of weight fragments. In this case, because the size of each of the plurality of files FL1-2 to FL4-2 is optimized or minimized, the DMA time between the memory device **210** and the accelerator **220** may be shortened. For example, when the size of a weight fragment is not considered, weight fragments with a relatively large size may be focused on a specific file or a specific memory. In this case, a DMA time for the specific file or the specific memory may increase. In contrast, according to the embodiment of FIGS. **13A** and **13B**, because weight fragments with a relatively large size (or weight fragments with a relatively small size) are distributed to the plurality of files FL1-2 to FL4-2, the length of each of the plurality of files FL1-2 to FL4-2 may become relatively short. In this case, the DMA time for each of the plurality of memories MEM1 to MEM4 may be shortened.

[0133] FIGS. **14A** and **14B** are diagrams for describing an operation of an aligning module of a host of FIG. **1**. Referring to FIGS. **1**, **14A**, and **14B**, the aligning module **120** of the host **100** may perform the aligning operations on the lightening weight data LWD or LWD1 to LWD4 based on the configuration information CONFIG and may generate a plurality of files FL1-3 to FL4-3. The lightening weight data LWD or LWD1 to LWD4 and the weight fragments W1a to W4d included in each of the lightening weight data LWD or LWD1 to LWD4 are described with reference to FIG. **4**, and thus, additional description will be omitted to avoid redundancy.

[0134] In the above embodiments, the aligning module **120** may distribute weight fragments corresponding to one weight data to different files, respectively. In contrast, in the embodiment of FIGS. **14A** and **14B**, the aligning module **120** of the host **100** may distribute “k” weight fragments (k being a natural number) corresponding to one weight data to “n” files (n being a natural number smaller than k). For example, the aligning module **120** may distribute the weight fragments W1a, W1b, W1c, and W1d of the first lightening weight data LWD1 to the first and second files FL1-3 and FL2-3, may distribute the weight fragments W2a, W2b, W2c, and W2d of the second lightening weight data LWD2 to the third and fourth files FL3-3 and FL4-3, may distribute the weight fragments W3a, W3b, W3c, and W3d of the third lightening weight data LWD3 to the first and second files FL1-3 and FL2-3, and may distribute the weight fragments W4a, W4b, W4c, and W4d of the fourth lightening weight data LWD4 to the third and fourth files FL3-3 and FL4-3.

[0135] In this case, the first file FL1-3 may include the 1a-th, 1d-th, 3a-th, and 3d-th weight fragments W1a, W1d, W3a, and W3d, the second file FL2-3 may include the 1b-th, 1c-th, 3b-th, and 3c-th weight fragments W1b, W1c, W3b, and W3c, the third file FL3-3 may include the 2a-th, 2b-th, 4c-th, and 4d-th weight fragments W2a, W2b, W4c, and W4d, and the fourth file FL4-3 may include the 2c-th, 2d-th, 4a-th, and 4b-th weight fragments W1c, W2d, W4a, and W4b.

[0136] A way to allocate weight fragments (e.g., a way to distribute weight fragments based on sizes of the weight fragments or a way to add padding data) is similar to that described above except that the number of weight fragments allocated to the same file from among weight fragments corresponding to the same lightening weight data is different from that described above, and thus, additional description will be omitted to avoid redundancy.

[0137] In at least one embodiment, the aligning module **120** of the host **100** may generate the first to fourth files FL1-3 to FL4-3 such that two weight fragments are read out from the first to fourth memories MEM1 to MEM4 in one DMA period. In this case, in each of the first to fourth memories MEM1 to MEM4, padding data may be added such that two weight fragments have the same length or size.

[0138] For example, as illustrated in FIG. **14A**, padding data may be added after the 1b-th and 1c-th weight fragments of the second file FL2-3 and after the 2a-th and 2b-th weight fragments W2a and W2b of the third file FL3-3 such that a first size of the 1a-th and 1d-th weight fragments W1a and W1d of the first file FL1-3, a second size of the 1b-th and 1c-th weight fragments W1b and

W1c of the second file FL2-3, a third size of the 2a-th and 2b-th weight fragments W2a and W2b of the third file FL3-3, and a fourth size of the 2c-th and 2d-th weight fragments W2c and W2d of the fourth file FL4-3 are identical to each other. Padding data may be added after the 4c-th and 4d-th weight fragments W4c and W4d of the third file FL3-3 such that a fifth size of the 3a-th and 3d-th weight fragments W3a and W3d of the first file FL1-3, a sixth size of the 3b-th and 3c-th weight fragments W3b and W3c of the second file FL2-3, a seventh size of the 4c-th and 4d-th weight fragments W4c and W4d of the third file FL3-3, and an eighth size of the 4a-th and 4b-th weight fragments W4a and W4b of the fourth file FL4-3 are identical to each other. In this case, the addition of the padding data may be minimized.

[0139] The first to fourth files FL1-3 to FL4-3 thus generated may be loaded into the first to fourth memories MEM1 to MEM4 and may be read out by the accelerator 220. For example, as illustrated in FIG. 14B, the first memory controller MCT1 may sequentially receive the 1a-th-1d-th, 3a-th, and 3d-th weight fragments W1a, W1d, W3a, and W3d from the first memory MEM1 through a first DMA operation. The second memory controller MCT2 may sequentially receive the weight 1b-th, 1c-th, 3b-th, and 3c-th weight fragments W1b, W1c, W3b, and W3c from the second memory MEM2 through a second DMA operation. The third memory controller MCT3 may sequentially receive the weight 2a-th, 2b-th, 4c-th, and 4d-th weight fragments W2a, W2b, W4c, and W4d from the third memory MEM3 through a third DMA operation. The fourth memory controller MCT4 may sequentially receive the 2c-th-2d-th, 4a-th, and 4b-th weight fragments W2c, W2d, W4a, and W4b from the fourth memory MEM4 through a fourth DMA operation.

[0140] At a first point in time t1, the accelerator 220 may generate the first lightening weight data LWD1 by using the 1a-th, 1b-th, 1c-th, and 1d-th weight fragments W1a, W1b, W1c, and W1d and may generate the second lightening weight data LWD2 by using the 2a-th, 2b-th, 2c-th, and 2d-th weight fragments W2a, W2b, W2c, and W2d. At a second point in time t2, the accelerator 220 may generate the third lightening weight data LWD3 by using the 3a-th, 3b-th, 3c-th, and 3d-th weight fragments W3a, W3b, W3c, and W3d and may generate the fourth lightening weight data LWD4 by using the 4a-th, 4b-th, 4c-th, and 4d-th weight fragments W4a, W4b, W4c, and W4d.

[0141] FIG. 15 is a diagram for describing an operation of an aligning module of a host of FIG. 1. Referring to FIGS. 1 and 15, the aligning module 120 of the host 100 may perform the aligning operations on the lightening weight data LWD or LWD1 to LWD4 based on the configuration information CONFIG and may generate a plurality of files FL1-4 to FL6-4. The lightening weight data LWD or LWD1 to LWD4 and the weight fragments W1a to W4d included in each of the lightening weight data LWD or LWD1 to LWD4 are described with reference to FIG. 4, and thus, additional description will be omitted to avoid redundancy.

[0142] In the present example, the memory device 210 and the accelerator 220 of the computing device 200 communicate with each other through the four channels CH1 to CH4. In this case, the aligning module 120 may generate the four files FL1 to FL4, FL1-1 to FL4-1, FL1-2 to FL4-2, or FL1-3 to FL4-3 to be stored in the memories MEM1 to MEM4 respectively corresponding to the four channels CH1 to CH4. However, the present disclosure is not limited thereto.

[0143] As an example, the memory device 210 and the accelerator 220 of the computing device 200 may communicate with each other through six channels. In this case, the aligning module 120 may generate sixth files FL1-4 to FL6-4. For example, the aligning module 120 may distribute the weight fragments W1a, W1b, W1c, and W1d of the first lightening weight data LWD1 to the first, second, third, and fourth files FL1-4, FL2-4, FL3-4, and FL4-4. The aligning module 120 may distribute the weight fragments W3a, W3b, W3c, and W3d of the third lightening weight data LWD3 to the first, second, third, and fourth files FL1-4, FL2-4, FL3-4, and FL4-4. The aligning module 120 may distribute the weight fragments W4a, W4b, W4c, and W4d of the fourth lightening weight data LWD4 to the first, second, third, and fourth files FL1-4, FL2-4, FL3-4, and FL4-4. The aligning module 120 may distribute the weight fragments W2a, W2b, W2c, and W2d of the second lightening weight data LWD2 to the fifth and sixth files FL5-4 and FL6-4.

[0144] In this case, the first file FL1-4 may include the 1a-th, 3a-th, and 4a-th weight fragments W1a, W3a, and W4a, the second file FL2-4 may include the 1b-th, 3b-th, and 4b-th weight fragments W1b, W3b, and W4b, the third file FL3-4 may include the 1c-th, 3c-th, and 4c-th weight fragments W1c, W3c, and W4c, the fourth file FL4-4 may include the 1d-th, 3d-th, and 4d-th weight fragments W1d, W3d, and W4d, the fifth file FL5-4 may include the 2a-th and 2c-th weight fragments W2a and W2c, and the sixth file FL6-4 may include the 2b-th and 2d-th weight fragments W2b and W2d.

[0145] The aligning module 120 may add padding data such that pieces of data to be accessed in one DMA period have the same sizes. A way to add padding data is described above, and thus, additional description will be omitted to avoid redundancy.

[0146] The first to sixth files FL1-4 to FL6-4 may be respectively loaded into six memories configured to communicate with the accelerator 220 through the six channels. The accelerator 220 may access the first to sixth files FL1-4 to FL6-4 from the six memories through the six channels independently or in parallel and may generate the first to fourth lightening weight data LWD1 to LWD4.

[0147] FIG. 16 is a diagram for describing an operation of an aligning module of a host of FIG. 1. Referring to FIGS. 1 and 16, the aligning module 120 of the host 100 may perform the aligning operations on the lightening weight data LWD or LWD1 to LWD4 based on the configuration information CONFIG and may generate a plurality of files FL1-5 to FL6-5. The lightening weight data LWD or LWD1 to LWD4 and the weight fragments W1a to W4d included in each of the lightening weight data LWD or LWD1 to LWD4 are described with reference to FIG. 4, and thus, additional description will be omitted to avoid redundancy.

[0148] In the embodiment of FIG. 16, the aligning module 120 may distribute weight fragments corresponding to one weight data to two files. For example, the aligning module 120 may distribute the weight fragments W1a, W1b, W1c, and W1d of the first lightening weight data LWD1 to the first and second files FL1-5 and FL2-5. The aligning module 120 may distribute the weight fragments W2a, W2b, W2c, and W2d of the second lightening weight data LWD2 to the third and fourth files FL3-5 and FL4-5. The aligning module 120 may distribute the weight fragments W3a, W3b, W3c, and W3d of the third lightening weight data LWD3 to the fifth and sixth files FL5-5 and FL6-5. The aligning module 120 may distribute the weight fragments W4a, W4b, W4c, and W4d of the fourth lightening weight data LWD4 to the first, second, third, and fourth files FL1-4, FL2-4, FL3-4, and FL4-4.

[0149] In this case, the first file FL1-5 may include the 1a-th, 1b-th, and 4a-th weight fragments W1a, W1b, and W4a, the second file FL2-5 may include the 1c-th, 1d-th, and 4b-th weight fragments W1c, W1d, and W4b, the third file FL3-5 may include the 2a-th, 2c-th, and 4c-th weight fragments W2a, W2c, and W4c, the fourth file FL4-5 may include the 2b-th, 2d-th, and 4d-th weight fragments W2b, W2d, and W4d, the fifth file FL5-5 may include the 3a-th and 3c-th weight fragments W3a and W3c, and the sixth file FL6-5 may include the 3b-th and 3d-th weight fragments W3b and W3d.

[0150] The first to sixth files FL1-4 to FL6-4 may be respectively loaded into six memories configured to communicate with the accelerator 220 through the six channels. The accelerator 220 may access the first to sixth files FL1-4 to FL6-4 from the six memories through the six channels independently or in parallel and may generate the first to fourth lightening weight data LWD1 to LWD4.

[0151] As described above, the host 100 according to at least one embodiment of the present disclosure may generate a plurality of files by performing the lightening operation and the aligning operation on a plurality of weight data WD based on the configuration information CONFIG of the computing device 200. The host 100 may load the plurality of files into the memory device 210 of the computing device 200. In this case, the accelerator 220 of the computing device 200 may access the plurality of files from the memory device 210 in parallel. This may mean that a time taken for

the accelerator **220** to obtain weight data is shortened.

[0152] In at least one embodiment, the host **100** may provide the computing device **200** with weight fragment information about weight fragments included in the plurality of files (e.g., with a start address and an offset of each of the weight fragments). The computing device **200** may read out the plurality of files from the memory device **210** based on the weight fragment information and may identify the weight fragments from the plurality of files. For example, the accelerator **220** may perform DMA setting for the plurality of memory controllers MCT**1** to MCT**4** based on the weight fragment information. The plurality of memory controllers MCT**1** to MCT**4** may perform the DMA operations on the plurality of memories MEM**1** to MEM**4** of the memory device **210** based on the DMA setting, and thus, the weight fragments may be identified.

[0153] In at least one embodiment, the host **100** may provide the weight fragment information to the computing device **200** through a separate communication line. Alternatively, the host **100** may load the weight fragment information into a given region of the memory device **210**. The accelerator **220** of the computing device **200** may access the given region of the memory device **210** to obtain the weight fragment information and may access the plurality of files based on the obtained weight fragment information.

[0154] The above embodiments are provided to describe the present disclosure easily, and the present disclosure is not limited thereto. For example, the configuration information CONFIG of the computing device **200** may include information about n-bit quantization and information about “k” channels. In this case, each of a plurality of lightening weight data generated by the lightening module **110** of the host **100** may be divided into “n” weight fragments, the sizes of which are different from each other. The aligning module **120** of the host **100** may generate a plurality of files by performing the aligning operation on the weight fragments. The plurality of files may be loaded into the memory device **210** of the computing device **200**, and the accelerator **220** of the computing device **200** may read out the plurality of files from the memory device **210**. In this case, the aligning module **120** may generate the plurality of files based on the configuration information CONFIG such that the accelerator **220** of the computing device **200** efficiently reads out the plurality of files from the memory device **210**. Accordingly, the performance of the artificial intelligence system **10** or the computing device **200** is improved.

[0155] FIG. **17** is a block diagram illustrating an artificial intelligence system according to at least one embodiment of the present disclosure. Referring to FIG. **17**, an artificial intelligence system **1000** may include a processor **1100**, a memory **1200**, a storage device **1300**, an artificial intelligence storage device **1400**, an artificial intelligence memory **1500**, a first accelerator **1610**, and a second accelerator **1620**. The processor **1100** may be configured to control all the operations of the artificial intelligence system **1000**. The processor **1100** may be configured to drive or execute various operating systems or various programs. The memory **1200** may be configured to be used as a system memory, a buffer memory, and/or a cache memory of the artificial intelligence system **1000**. The storage device **1300** may be used as a high-capacity storage medium of the artificial intelligence system **1000**.

[0156] As an example, the storage device **1300** may be configured to store information or data associated with various operating systems or various programs which are driven in the artificial intelligence system **1000**. The data or information stored in the storage device **1300** may be loaded into the memory **1200**, and the processor **1100** may drive or execute the operating systems or programs based on the information or data loaded into the memory **1200**.

[0157] The artificial intelligence storage device **1400** may be configured to store various information, an artificial intelligence model, or weight data for an artificial intelligence computation, learning, or inference. For example, the AI artificial intelligence storage device **1400** may be and/or include the storage device **300** of FIG. **1**. In at least one embodiment, the processor **1100** may generate a plurality of files by performing the lightening operation and the aligning operation on the information, the artificial intelligence model, or the weight data which are stored

in the artificial intelligence storage device **1400** for the artificial intelligence computation, the learning, or the inference. The plurality of files may be stored in the artificial intelligence memory **1500**. In at least one embodiment, the processor **1100** may perform the lightening operation and the aligning operation based on the method described with reference to FIGS. **1** to **16**. For example, the processor **1100** may be and/or include the host **100** of FIG. **1**.

[0158] The first and second accelerators **1610** and **1620** may read out the plurality of files from the artificial intelligence memory **1500** and may generate various information, an artificial intelligence model, or weight data for an artificial intelligence computation, learning, or inference based on the plurality of files. The first and second accelerators **1610** and **1620** may perform the artificial intelligence computation, the learning, or the inference based on the generated information. For example, the first and second accelerators **1610** and **1620** may each be or include an accelerator **220** of FIG. **1**.

[0159] In at least one embodiment, the first and second accelerators **1610** and **1620** may share the artificial intelligence memory **1500**. For example, the first and second accelerators **1610** and **1620** may access the artificial intelligence memory **1500** through the same channels. Alternatively, the first and second accelerators **1610** and **1620** may access different regions of the artificial intelligence memory **1500**. For example, the first accelerator **1610** may access a first region of the artificial intelligence memory **1500** through first channels, and the second accelerator **1620** may access a second region of the artificial intelligence memory **1500** through second channels. Otherwise, the artificial intelligence memory **1500** may be the same as and/or substantially similar to the memory device **210** of FIG. **1**.

[0160] In at least one embodiment, the processor **1100** may perform the lightening operation and the aligning operation based on configuration information of each of the first and second accelerators **1610** and **1620** such that each of the first and second accelerators **1610** and **1620** efficiently accesses the artificial intelligence memory **1500**.

[0161] In the embodiment of FIG. **17**, individual components (e.g., the artificial intelligence storage device **1400** and the artificial intelligence memory **1500**) for an artificial intelligence computation, learning, or inference are illustrated, but the present disclosure is not limited thereto. The artificial intelligence storage device **1400** and/or the artificial intelligence memory **1500** may be omitted depending on the implementation of the artificial intelligence system; in this case, data, information, or weight data for an artificial intelligence computation, learning, or inference may be stored in the storage device **1300**, and a plurality of files generated through the lightening operation or the aligning operation of the processor **1100** may be stored in the memory **1200**. Accordingly, the first and second accelerators **1610** and **1620** may access the plurality files from the memory **1200**.

[0162] According to the present disclosure, a host configured to control a computing device may generate a plurality of files based on configuration information of the computing device such that an accelerator of the computing device efficiently accesses an artificial intelligence model. The plurality of files are respectively loaded into a plurality of memories of the computing device. In this case, the accelerator of the computing device may read out the plurality of files by performing direct memory access (DMA) operations from the plurality of memories in parallel. Accordingly, an operation method of a host controlling a computing device with improved performance and an operation method of an artificial intelligence system including the computing device and the host are provided.

[0163] While the present disclosure has been described with reference to some embodiments thereof, it will be apparent to those of ordinary skill in the art that various changes and modifications may be made thereto without departing from the spirit and scope of the present disclosure as set forth in the following claims.

Claims

1. An operation method of a host configured to control a computing device performing an artificial intelligence computation, the method comprising: receiving configuration information from the computation device; generating a plurality of lightening weight data by performing lightening on weight data based on the configuration information; generating a plurality of files by performing an aligning operation on the plurality of lightening weight data based on the configuration information; and loading the plurality of files into a memory device of the computing device, wherein the configuration information includes channel information about a plurality of channels between the memory device of the computing device and an accelerator of the computing device, and wherein a number of the plurality of files generated is equal to a number of available channels of the plurality of channels.
2. The method of claim 1, wherein the configuration information includes quantization information, and wherein the generating of the plurality of lightening weight data includes generating a plurality of unit data by performing quantization on the weight data based on the quantization information; and generating the plurality of lightening weight data by performing a compression operation on the plurality of unit data.
3. The method of claim 2, wherein the generating the plurality of lightening weight data includes generating the plurality of lightening weight data such that each of the plurality of lightening weight data includes a plurality of weight fragments and such that a number of the plurality of weight fragments included in each of the plurality of lightening weight data is based on the quantization information.
4. The method of claim 3, wherein the generating of the plurality of files by performing the aligning operation on the plurality of lightening weight data based on the configuration information includes: distributing the plurality of weight fragments included in each of the plurality of lightening weight data to the plurality of files based on the configuration information.
5. The method of claim 4, further comprising: providing weight fragment information about the plurality of files to the computing device, wherein the weight fragment information includes information about a start address and an offset of each of the plurality of weight fragments included in each of the plurality of files.
6. The method of claim 3, wherein the generating the plurality of files includes generating the plurality of files such that a first file among the plurality of files includes a first part of a plurality of first weight fragments included in first lightening weight data among the plurality of lightening weight data, and such that a second file among the plurality of files includes a second part of the plurality of first weight fragments included in the first lightening weight data among the plurality of lightening weight data.
7. The method of claim 6, wherein the loading the plurality of files into the memory device of the computing device includes: loading the first file into a first memory communicating with the accelerator through a first channel from among a plurality of memories of the memory device, and loading the second file into a second memory communicating with the accelerator through a second channel from among the plurality of memories of the memory device.
8. The method of claim 7, wherein the loading the plurality of files into the memory device of the computing device includes loading the plurality of files into the memory device such that a start address of the first part included in the first file loaded into the first memory is identical to a start address of the second part included in the second file loaded into the second memory.
9. The method of claim 6, wherein the generating the plurality of files includes generating the plurality of files such that a size of the first part is larger than a size of the second part, and such that the second file further includes first padding data added to the second part.
10. The method of claim 1, wherein the generating the plurality of files includes generating the

plurality of files such that sizes of the plurality of files are identical to each other.

11. The method of claim 1, wherein the generating the plurality of lightening weight data includes: reading out the weight data from a storage device.

12. An operation method of a host configured to control a computing device performing an artificial intelligence computation, the method comprising: receiving configuration information from the computation device; generating first lightening weight data including first weight fragments and second lightening weight data including second weight fragments by performing a lightening operation on first weight data and second weight data based on the configuration information; generating a plurality of files by performing an aligning operation on the first weight fragments and the second weight fragments; and loading the plurality of files into a plurality of memories of the computing device, wherein the configuration information includes channel information about a plurality of channels between the plurality of memories of the computing device and an accelerator of the computing device, wherein a number of the plurality of files generated is equal to a number of the plurality of channels, wherein the first weight fragments are distributed to the plurality of files, and wherein the second weight fragments are distributed to the plurality of files.

13. The method of claim 12, wherein the configuration information includes quantization information, and wherein the generating of the first lightening weight data including the first weight fragments and the second lightening weight data including the second weight fragments by performing the lightening operation on the first weight data and the second weight data based on the configuration information includes generating a plurality of first unit data and a plurality of second unit data by performing quantization operations on the first weight data and the second weight data, respectively, based on the quantization information, and generating the first lightening weight data including the first weight fragments and the second lightening weight data including the second weight fragments by performing a compression operation on the plurality of first unit data and a compression operation on the plurality of second unit data.

14. The method of claim 12, wherein generating first lightening weight data including first weight fragments and second lightening weight data including second weight fragments includes reading out the first weight data and the second weight data from a storage device.

15. The method of claim 12, further comprising: providing weight fragment information about the plurality of files to the computing device, wherein the weight fragment information includes information about an offset and a start address corresponding to each of the first weight fragments and the second weight fragments, and wherein the loading the plurality of files into the plurality of memories of the computing device includes loading the plurality of files into the plurality of memories such that the first weight fragments and the second weight fragments are stored in the plurality of memories based on the start address and the offset corresponding thereto.

16. An operation method of an artificial intelligence system which includes a computing device configured to perform an artificial intelligence computation and a host configured to communicate with the computing device, the method comprising: receiving, by the host, configuration information of the computation device; generating, by the host, a plurality of files by performing a lightening operation and an aligning operation on weight data based on the configuration information; loading, by the host, the plurality of files into a plurality of memories of the computing device; reading out, by the computing device, the plurality of files by performing direct memory access (DMA) operations on the plurality of memories in parallel; and performing, by the computing device, the artificial intelligence computation based on the plurality of files, wherein the configuration information includes channel information about a plurality of channels between the plurality of memories of the computing device and an accelerator of the computing device, and wherein the generating, by the host, the plurality of files includes generating the plurality of files such that a number of the plurality of files is equal to a number of the plurality of channels.

17. The method of claim 16, wherein the configuration information includes quantization

information, and wherein the lightening operation includes generating, by the host, a plurality of unit data by performing a quantization operation on the weight data based on the quantization information, and generating, by the host, a plurality of lightening weight data by performing a compression operation on the plurality of unit data.

18. The method of claim 17, wherein the generating, by the host, the plurality of files includes generating the plurality of files such that each of the plurality of lightening weight data includes a plurality of weight fragments, and such that the number of the plurality of weight fragments included in each of the plurality of lightening weight data corresponds to the quantization information, and wherein the aligning operation includes generating, by the host, the plurality of files by distributing the plurality of weight fragments included in each of the plurality of lightening weight data to the plurality of files.

19. The method of claim 18, wherein the reading out, by the computing device, of the plurality of files includes: reading out, by the computing device, first weight fragments, among the plurality of weight fragments, from the plurality of memories, respectively, by performing the DMA operations on the plurality of memories in parallel, during a first time interval; generating first weight data based on a first part of the first weight fragments; and buffering a second part of the first weight fragments.

20. The method of claim 19, wherein the reading out, by the computing device, of the plurality of files, by performing the DMA operations on the plurality of memories in parallel further includes: reading out, by the computing device, second weight fragments, among the plurality of weight fragments, from the plurality of memories, respectively, by performing the DMA operations on the plurality of memories in parallel, during a second time interval; generating second weight data based on a third part of the second weight fragments and the buffered second part; and buffering a fourth part of the second weight fragments.
