

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258693

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

KACHARE; Ramdas P. et al.

STORAGE DEVICE-ASSISTED LIVE VIRTUAL MACHINE MIGRATION

Abstract

A storage device is disclosed. The storage device may include at least one controller for a virtual machine (VM) that is on a source host. Storage in the storage device may store data for the VM. A second storage may store a storage state for the VM. A storage device controller may process at least one read request received from the controller for the VM using the first storage and at least one write request received from the controller for the VM using the first storage. A VM migration state monitor and capture module may assist in the migration of the VM from the source host to a destination host.

Inventors: KACHARE; Ramdas P. (Pleasanton, CA), PINTO; Oscar P. (San Jose, CA), KI; Yang Seok (Palo Alto, CA)

Applicant: Samsung Electronics Co., Ltd. (Suwon-si, KR)

Family ID: 77365222

Appl. No.: 19/060619

Filed: February 21, 2025

Related U.S. Application Data

parent US continuation 17006773 20200828 parent-grant-document US 12254335 child US 19060619

us-provisional-application US 63040516 20200617

us-provisional-application US 62982057 20200226

Publication Classification

Int. Cl.: G06F9/455 (20180101); G06F3/06 (20060101); G06F12/0862 (20160101)

U.S. Cl.:

CPC **G06F9/45558** (20130101); **G06F3/0607** (20130101); **G06F3/0634** (20130101); **G06F3/0647** (20130101); **G06F3/0658** (20130101); **G06F3/0664** (20130101); **G06F12/0862** (20130101); G06F2009/4557 (20130101); G06F2009/45579 (20130101); G06F2009/45591 (20130101)

Background/Summary

RELATED APPLICATION DATA [0001] This application is a continuation of U.S. patent application Ser. No. 17/006,773, filed Aug. 28, 2020, which claims the benefit of U.S. Provisional Patent Application Ser. No. 63/040,516, filed Jun. 17, 2020, and U.S. Provisional Patent Application Ser. No. 62/982,057, filed Feb. 26, 2020; U.S. Provisional Patent Application Ser. No. 63/040,516, filed Jun. 17, 2020, is incorporated by reference herein for all purposes.

FIELD

[0002] The inventive concepts relate generally to storage devices, and more particularly to a storage device supporting live migration of a virtual machine.

BACKGROUND

[0003] The amount of data being generated, analyzed, and consumed continues to grow. Among the technologies being employed for efficient utilization of the data centers is virtualization. The virtualization technologies enable a physical server to run multiple Virtual Machines (VM) for the user applications on a single physical server. This technique allows data centers to support large number of user applications with existing resources and reduces underutilization of the physical resources.

[0004] Data centers may move (i.e., migrate) a virtual machine from one server to another. Migrating VMs may use resources of the source and destination machine, and may involve temporarily stopping the VM being migrated. Thus, it is desirable for VM migration to be as efficient as possible.

[0005] A need remains to improve the migration of VMs from one server to another.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0006] FIG. 1 shows a machine including a storage device equipped to assist in live migration, according to an embodiment of the inventive concept.

[0007] FIG. 2 shows additional details of the machine of FIG. 1, according to embodiments of the inventive concept.

[0008] FIG. 3 shows the machine and storage device of FIG. 1 supporting virtual machines (VMs), according to embodiments of the inventive concept.

[0009] FIG. 4 shows details of the storage device of FIG. 1, according to embodiments of the inventive concept.

[0010] FIG. 5 shows a virtual machine being migrated from one machine to another, according to embodiments of the inventive concept.

[0011] FIG. 6 shows details of state information that may be maintained by the storage device of FIG. 1, according to embodiments of the inventive concept.

[0012] FIG. 7 shows the VM Migration State Monitor and Capture Module of FIG. 4 determining whether a threshold number of logical block addresses (LBAs) have been written, according to embodiments of the inventive concept.

[0013] FIG. 8 shows the VM Migration State Monitor and Capture Module of FIG. 4 generating a log page for the state of the storage device of FIG. 1, according to embodiments of the inventive concept.

[0014] FIGS. 9A-9B show a flowchart of an example procedure for the storage device of FIG. 1 to support live migration of a VM from the machine of FIG. 1, according to embodiments of the inventive concept.

[0015] FIG. 10 shows a flowchart of an example procedure for the VM Migration State Monitor and Capture Module of FIG. 4 to determine whether a threshold number of LBAs have been written, according to an embodiment of the inventive concept.

[0016] FIG. 11 shows a flowchart of an example procedure for the VM Migration State Monitor and Capture Module of FIG. 4 to generate and deliver a report of the storage state of FIG. 4 for the VM, according to an embodiment of the inventive concept.

[0017] FIGS. 12A-12B show a flowchart of an example procedure for the storage device of FIG. 1 to support live migration of a VM to the machine of FIG. 1, according to embodiments of the inventive concept.

DETAILED DESCRIPTION

[0018] Reference will now be made in detail to embodiments of the inventive concept, examples of which are illustrated in the accompanying drawings. In the following detailed description, numerous specific details are set forth to enable a thorough understanding of the inventive concept. It should be understood, however, that persons having ordinary skill in the art may practice the inventive concept without these specific details. In other instances, well-known methods, procedures, components, circuits, and networks have not been described in detail so as not to unnecessarily obscure aspects of the embodiments.

[0019] It will be understood that, although the terms first, second, etc. may be used herein to describe various elements, these elements should not be limited by these terms. These terms are only used to distinguish one element from another. For example, a first module could be termed a second module, and, similarly, a second module could be termed a first module, without departing from the scope of the inventive concept.

[0020] The terminology used in the description of the inventive concept herein is for the purpose of describing particular embodiments only and is not intended to be limiting of the inventive concept. As used in the description of the inventive concept and the appended claims, the singular forms “a”, “an”, and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will also be understood that the term “and/or” as used herein refers to and encompasses any and all possible combinations of one or more of the associated listed items. It will be further understood that the terms “comprises” and/or “comprising,” when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof. The components and features of the drawings are not necessarily drawn to scale.

[0021] As the amount of data being generated, analyzed, and consumed is growing, the data center scale is growing to keep pace with the data growth. New business models of cloud based IT resources are emerging. The cloud computing is enabling users to have more economical IT infrastructure using on-demand, pay-as-you-go, and use-what-you-need resources providing significant cost benefits to the end users. Due to these factors, the data centers are becoming complex, large, and involve outlays of capital expenditures (capex) and operational expenditures (opex). Hence the efficient utilization of the data center resources such as servers has become important.

[0022] Among the many technologies being developed and employed for efficient utilization of the data centers is virtualization. The virtualization technologies enable a physical server to run multiple Virtual Machines (VM) for the user applications on a single physical server. This

technique allows data centers to support large number of user applications with existing resources and reduces underutilization of the physical resources.

[0023] The virtualization Operating Systems (OS), also known as a hypervisor, may run on the physical servers and facilitate VM operations and management. Hypervisors enable launching of the VMs, resource allocations to the VMs, isolation and fairness between multiple VMs, as well as migration of VMs from one server to another. Each VM needs four primary resources—namely 1) compute 2) system memory, 3) network connectivity, and 4) storage—for smooth and correct operations. A hypervisor may provision, manage, and allocate these resources to the VMs running on that server. In a typical data center, thousands of physical servers may run a very large number of VMs. Thus, load balancing of the VMs on the given set or cluster of servers may also be important.

[0024] At times, maintenance and upgrades are needed to be performed to these servers. Sometimes, the requirement of the VM may be upgraded, be it by resource requirements or user requirements. It is also possible that some servers may be running too few VMs and it may make sense to consolidate these VMs on another server for cost savings. All of these reasons lead to a need to have VM migration capability in the servers and the hypervisors.

[0025] The VM migration should be as non-intrusive and non-disruptive as possible. Ideally, the applications running on the VMs should remain agnostic to the fact that the underlying VMs are being migrated to different servers. This process is also known as live migration. Hence, it is desirable to reduce and minimize the time taken to move a VM from the source server to the destination server. When a hypervisor migrates a VM, it essentially copies the state of the VM from the source server to the destination server. The state of a VM consists of processor state, memory state, network state, and storage state.

[0026] The VM migration is typically done in three phases: 1) preparation, 2) data copy (brownout), and 3) delta copy (blackout), as further described below. In the preparation phase, various resources needed for the candidate VM are allocated or provisioned on the destination server. This negotiation happens between the source and destination servers. Once the resources are reserved and allocated for the VM on the destination server, the second phase starts. In the data copy phase, also known as brownout phase, the source hypervisor starts moving VM data to the destination server. In this phase both the VM and the hypervisor actively access the VM data on the source Solid State Drive (SSD). It is likely that some of the data being moved to the destination may be modified by the VM during this phase. Once the whole data copy is done, live migration enters the delta copy phase, also called as blackout. In this phase, the source hypervisor suspends or pauses the VM being migrated. Once the VM is frozen, that VM cannot change any data or state of its resources. At that point, the hypervisor copies over the delta changes compared to the copy operation it performed during second phase. Once the delta changes are copied over, the destination hypervisor launches the VM on the destination server, completing the live migration process.

[0027] The storage resource for a VM may be provisioned from a direct attached SSD or remotely from disaggregated storage. During the data copy operations mentioned above, the source server that hosts the VM may allocate compute, memory, and in some cases storage resources to keep track of the delta changes for the third phase. The state may be captured as a bitmap of individual or a range of LBAs. Similarly, the destination server may allocate similar resources for the transfer. These resources may add additional cost to the solution and inadvertently add latencies constraints on existing guests on these servers besides other limitations.

[0028] Embodiments of the inventive concept may include a system and method of using SSD assistance to perform fast and efficient live VM migration from a source server to the destination server. The VM migrations are done for various reasons such as server consolidation, load balancing etc. As part of the VM migration process, the data and state of various VM resources such as CPU, memory, storage, and networking may be replicated on the destination sever. During the state transfer and relaunch of the VM on the destination, the VM may be paused. It is desirable

that the pause of the VM is minimal so that the applications running on the VM are not affected. Embodiments of the inventive concept may use the SSD controller to assist in the movement of storage data to the destination server in an efficient manner.

[0029] When the hypervisor makes the decision to migrate the VM, the source hypervisor (that is, the hypervisor on the source server) may alert the SSD device (on the source server) of that decision. The SSD controller may then start tracking the storage state changes as well as data changes. The SSD controller may create dirty state and LBA ranges logs that minimize the pause time for the VM to transfer the storage data and state. The SSD controller may also adjust the data prefetch and write data cache policies during the VM migration active period so that data belonging to that VM may be accessed quickly and efficiently. The SSD controller may also keep track of the changes from the last snapshot and the delta may be then used by the source hypervisor in a streaming fashion. Embodiments of the inventive concept described herein may also be used in other storage use cases such as snap-shot, archival, and backup features to achieve efficient execution as well.

[0030] System Architecture A source server may host multiple VMs. Each VM may be assigned to a Physical Function and Non-Volatile Memory Express (NVMe) controller in an attached SSD. (While NVMe controllers are specifically discussed, embodiments of the inventive concept may use other protocols for communication, and should not be considered limited to NVMe controllers.) A server may have multiple SSDs attached to it and each SSD may contain one or more NVMe controllers. From this pool of NVMe controllers, also known as storage controllers, one or more controllers may be assigned to each VM hosted by the hypervisor virtualization Operating System (OS) running on the server. Each VM may have its own guest OS and system software stack and applications. Each VM may have access to the server resources such as central processing unit (CPU) bandwidth, system memory, network bandwidth, and storage bandwidth and capacity. These resources may be managed, allocated, and assigned by the hypervisor to each VM hosted on that server.

[0031] The destination server may have a setup and configuration similar to the source server. For various reasons such as load balancing, server consolidation, etc., one or more VMs from the source server may be migrated to the destination server.

VM Storage State

[0032] Each VM running on a hypervisor may have various state vectors at any given time. The VM state may have multiple components including processor state, applications execution status, memory state, network state and storage state. Regarding storage state, the VM state may indicate VM activities related to persistent storage of data. Each VM may have a storage device it is attached to through a pass-through mechanism via the hypervisor.

[0033] The storage device state in turn may have three primary components: 1) PCIe transport connection state, 2) NVMe Controller state, and 3) Flash Translation Layer (FTL) state.

[0034] The PCIe interface present between the source server and the SSD storage device may have certain configurations and parameters set by the source hypervisor or the VM PCIe driver. Some examples of the PCIe configuration state include: Various PCIe capability settings such as Peripheral Component Interconnect Express (PCIe) Express, message signaled interrupts (MSI and MSI-X), advanced error reporting (AER), power management, etc.; and Configuration space settings such as interrupt, max payload size, etc.

[0035] During the lifetime of the VM, the VM NVMe driver may perform various NVMe configurations and settings. The NVMe state may consist of NVMe registers exposed through the PCIe BAR window. The NVMe state may also consist of various feature settings and configurations done by the NVMe driver. Some examples of NVMe level storage device state include: memory page size, doorbell stride, timeout values, submission queue (SQ) arbitration settings, etc.; Feature settings such as power management, number of queues, asynchronous event notification, interrupt coalescing, atomicity, etc.; namespaces created, security settings, etc.; and

log pages created so far, etc.

[0036] The FTL state may correspond to a state of logical to physical address mapping, flash media block usage, write/read cache state and settings for the associated NVMe Controller and in turn associated VM. In some embodiments of the inventive concept, VM migration may reproduce the FTL state on the destination SSD as closely as possible to the source SSD. In other embodiments of the inventive concept, replication of the FTL state may not be performed.

[0037] It is also possible for the SSD to keep track of Input/Output (I/O) activity state. That is to say, the SSD controller may record VM data writes in terms of LBA ranges. In some embodiments of the inventive concept, the SSD controller may record write LBA ranges below certain threshold in terms of number of unique LBAs written. Thus, the SSD controller may track the LBA ranges written by the VM and flag indicating whether the number of LBAs written exceeds some threshold (which may vary depending on the VM).

VM Migration Optimized SSD Architecture

[0038] SSDs may use flash media for persistent data storage. An SSD controller may consist of the following major modules: 1) host interface, 2) Flash Translation layer (FTL), and 3) flash channels. The SSD controller architecture may enable the performance of the flash media to be made accessible to the host.

[0039] The host interface essentially provides a logical view of the persistent storage to the attached host. The flash storage capacity may be advertised in units of logical blocks or sectors, or it may be advertised as object storage. The host may issue I/O commands using logical block addresses (LBA) in case of block storage, or the host may also use object storage protocols such as Key Value (KV). The host interface may implement PCIe as a transport protocol to carry host storage protocol such as NVMe. The host may issue storage I/O and administration commands using command submission queues (SQ). The host interface may perform SQ arbitration and selects a command for execution. The host interface may perform data direct memory access (DMA) operations to/from host memory as necessary. After the command execution the host interface may send a completion entry (CE) to the host. The host interface may also send an interrupt such as MSI or MSI-X to the host to indicate that one or more commands are executed.

[0040] Once the host interface selects a command for execution, the host interface may fetch the command and interpret the command issued by the host. The host interface may extract appropriate fields such as Logical Block address (LBA), data transfer length, and source/destination address for the data. The host interface may provide these parameters to the Flash Translation Layer (FTL) for further execution of the command. The FTL module essentially converts the logical address to the actual physical address where data may reside in the flash media. The FTL layer may maintain a Logical-to-Physical address translation table and performs a look-up to get the physical address. In addition to the address translation, FTL may perform other functions necessary for the flash media management.

[0041] Flash media may have certain read/write and endurance characteristics that should be managed. The SSD controller may manage the flash media so that it may be used to provide reliable, high performance, and cost effective host data storage. There are a number of internal house-keeping operations that a SSD controller may perform for normal and healthy operation of a SSD. Some examples of internal or background operations are: recycling; garbage collection; FTL data structure updates; various stats and logs updates; and cache offloads, among other possibilities.

[0042] The FTL layer may provide the physical address and data in case of write commands to the flash channels. The flash channels may be used to perform all the actual data read or write operations to the flash media. The flash channels may implement specific bus protocols needed to access the media and they may also implement other important functions such as error detection and correction.

[0043] Embodiments of the inventive concept may implement a number of features to assist with

the live VM migration process. When the source server host decides to move a VM to another destination server, the source server may inform that decision to the PCIe physical function (PF)/NVMe Controller associated with that VM. The hypervisor may set a flag in the appropriate SSD and NVMe controller. The host may use NVMe Set Feature command to set a “VM_MIGRATION_ACTIVE” flag in the host interface. The host may also use custom, vendor-defined NVMe command to set VM_MIGRATION_ACTIVE flag in the SSD. These commands are just examples, and any other such mechanisms may be used by the host to set such a flag in the SSD.

[0044] While the above example focuses on the NVMe protocol, similar flags, operations, commands, and/or functionalities may be available using other protocols: no specific description of commands, operations, functions, and/or features relating to NVMe herein should be considered limited to NVMe, and comparable commands, operations, functions, and/or features of other protocols may be used instead.

[0045] In some embodiments of the inventive concept, the VM to be migrated may still be active (live) and the SSD may be fully owned by the VM (pass-through). In such embodiments of the inventive concept, the NVMe commands may be sent out through the management interface if the active path is not available to the hypervisor.

[0046] Once the VM_MIGRATION_ACTIVE flag is set in the SSD, the SSD may start monitoring all the I/O activities and state of that PF/NVMe controller. A VM migration state monitor and capture module may create a log or report of the VM storage state described earlier. If the host changes any state parameters during the migration phase, the SSD may update the log/report appropriately. At some point, the source hypervisor may read the state log or report from the source SSD using NVMe Get Log command. The source hypervisor may also use custom, vendor-defined NVMe command to read the storage state of the VM being migrated. The above commands are just examples, and other mechanisms are possible.

[0047] The source hypervisor may read the storage state log page and then send that log or report to the destination server. At the destination server, the destination hypervisor may use a NVMe command Set Feature to provide the Storage State log or report to the appropriate SSD and NVMe controller to be used for the migrated VM. The SSD controller may use the storage state log received from the host and may plug the settings and configurations into appropriate modules such as PCIe, NVMe controller, and FTL. Using this method, the storage state of the destination SSD/NVMe controller may be ready or very close to the ready state for the VM migration. This feature helps reduce the time needed to stop the VM on the source server and relaunch it on the destination server. Hence the downtime aka blackout time during live migration is reduced.

[0048] As mentioned above, once the VM_MIGRATION_ACTIVE flag is set, the SSD controller may keep track of the I/O operations being performed by the VM. The SSD controller may create and maintain a dirty log of the LBA ranges that are modified by the VM. After setting VM_MIGRATION_ACTIVE flag, the source hypervisor may move all the flash data blocks (LBAs, namespaces, etc.) to the destination server. The source hypervisor may perform a sequential read of all the data belonging to the VM and send that data to the destination hypervisor. The destination hypervisor may then write that data into the associated SSD and NVMe controller on the destination server. In another embodiment of the inventive concept, the destination hypervisor may have already prepared the storage at the destination server with the last known backed up image or the original VM image and created checksums of LBAs by individual count or by ranges. In that case, the destination hypervisor may compare the old and new checksums and only copy the LBAs that do not match.

[0049] In one embodiment of the inventive concept, the SSD may keep track of the actual LBA ranges written by the VM during its lifetime, as part of the storage state for each VM. This information may be part of the storage state log or report page. Using this information, the source hypervisor may not read out the whole flash capacity allocated to the VM. Instead, the source

hypervisor may just read the actual LBA ranges written by the VM so that the amount of data moved may be reduced.

[0050] In some embodiments of the inventive concept, there may be a threshold beyond which the source hypervisor may just move the whole capacity as a sequential read workload. For example, consider a situation where approximately 100 GB of storage capacity was allocated to a VM. If the VM has only written approximately 4 GB of LBAs before the source hypervisor decided to migrate the VM, the source hypervisor may move just the actual LBAs to the destination hypervisor, which may be less than having to move the entire (approximately 100 GB) allocated to the VM. But by using a threshold, the burden on the source hypervisor may be reduced. If the VM has written, say, approximately 98 GB of LBAs before the source hypervisor decided to migrate the VM, the LBAs written information is not as helpful, as pretty much the whole allocated capacity needs to be moved to the destination server or destination SSD. Thus, if the source hypervisor detects that the threshold flag is set (indicating that the LBAs written is greater than some threshold), the source hypervisor may copy the whole data capacity to the destination server.

[0051] As mentioned earlier, once the source hypervisor sets the VM_MIGRATION_ACTIVE flag in the SSD controller, the SSD controller may create a dirty log for the LBAs written by the VM from that point in time onwards. The dirty log may be used to copy the modified data after initial whole data move. While the source hypervisor is moving data to the destination server, the VM may still be active in a phase sometimes referred as brownout. This is the phase during which both the source hypervisor and the VMs may actively access the SSD. Any LBAs modified during this phase may be copied over to the destination server later, before the VM is launched on the destination server. At the end of the brownout phase, the source hypervisor may stop or pause the VM on the source server, copy over any delta changes in state and data to the destination server, and finally launch the VM on the destination server. This final phase is also sometimes referred as blackout phase.

[0052] The objective is to minimize the length or the time spent in the blackout phase. The blackout time is the time during which VM is disrupted. Embodiments of the inventive concept may help fast completion of the blackout phase using assistance from SSD. After the source hypervisor stops the VM, no more state changes or LBA data changes happen in the source SSD. At that point, the source hypervisor may read the dirty log from the SSD to discover what LBAs and state parameters were modified during brownout phase. The source hypervisor may then copy the delta state changes and/or modified (or delta or dirty) LBAs to the destination. Once the delta state and LBA data is copied over to the destination, the VM may be launched on the destination server. At that point, live migration of the VM is successfully completed.

[0053] In another embodiment of the inventive concept, the SSD controller may take other steps to assist in the live VM migration process. Once the VM_MIGRATION_ACTIVE flag is set, the SSD controller may understand that the live VM migration is in progress, and the source hypervisor is going to copy that VM data to the destination. The SSD controller may start prefetching some of the sequential data or LBA ranges of interest from the flash media. In some embodiments of the inventive concept, data stored in the flash media may be encrypted: data decryption may take additional time. Prefetching of such encrypted data may result in data decryption in advance of the data being requested for migration. Such data prefetch policy may result in quicker completion of the hypervisor data read commands. Such reduction in the latency of the source hypervisor read commands may help reduce overall brownout phase time duration (with less data being modified during the brownout phase, the blackout phase may need less time to copy the reduced modified data).

[0054] In yet another embodiment of the inventive concept, once the VM_MIGRATION_ACTIVE flag is set, the SSD controller may provide higher priority to the associated VM data for read and write caching. That is to say, the SSD controller may cache the data being accessed by that VM on priority. Such caching may expedite command completions for both the VM being migrated as well

as the source hypervisor.

[0055] In yet another embodiment of the inventive concept, the source SSD may keep track of all the changes with respect to the last known image. The destination server SSD may be in sync with this last known image. Here both the SSDs understand the image state definitions. Once the source server is ready to transfer the VM, the source server may instruct the source and destination SSDs to get ready for stream transfers. In stream transfer mode, the SSD may prepare data to transfer and once started the host server may facilitate the transfer of the stream of data from source to destination SSD. The source and destination server may transfer data as is without decoding it.

Migration Assist NVMe Commands

[0056] The source hypervisor may use NVMe command extensions described below for configuring live migration assistance features in the SSD controller. These NVMe commands use standard format and protocol with some fields either repurposed or assigned new syntax. New custom, vendor-defined NVMe commands may also be used for this purpose. It may also be possible for the source hypervisor and the SSD controller to use an entirely different communication method to achieve the same functionality proposed. The following sections describe some examples of possible NVMe command extensions that may be used to facilitate live VM migration assistance from SSD.

[0057] NVMe Set Features—VM MIGRATION ACTIVE command may be used by the source to set certain NVME controller features in the SSD. Some examples of the existing features set by the source may include: command submission queue arbitration parameters; power management settings; temperature thresholds; number of queues; and interrupt coalescing configuration.

[0058] The source hypervisor may also set the following parameters:

VM_MIGRATION_ACTIVE; LBA_WRITTEN_THRESHOLD; and
LBA_WRITTEN_THRESHOLD_EXCEEDED.

[0059] VM_MIGRATION_ACTIVE may be a flag set by the source hypervisor. This flag may indicate that the hypervisor is in the process of migrating the VM associated that PF/NVMe controller. When this flag is set, the SSD controller may prepare certain storage state logs and reports, described in earlier sections, pertinent to that PF/NVMe controller. The SSD controller may also start creating dirty logs for the I/O activity by the VM. The SSD controller may also employ different set of data prefetch and data caching policies when this flag is set.

[0060] LBA_WRITTEN_THRESHOLD may be used by the SSD controller to keep track of the actual LBA ranges that are written by the VM during its lifespan. If the threshold value is set to 0, all the LBA ranges written by the VM may be recorded in a log. For non-zero values of the threshold, LBA ranges written by the host are recorded only up to the threshold size set by the host. Beyond the threshold size, the SSD controller may set the

LBA_WRITTEN_THRESHOLD_EXCEEDED flag for the source to read. During VM migration, if the source hypervisor finds that the LBA_WRITTEN_THRESHOLD_EXCEEDED flag was set, the source hypervisor may just move the whole capacity allocated to that VM to the destination server. If the LBA_WRITTEN_THRESHOLD_EXCEEDED flag was not set, the source hypervisor may only move the recorded LBA ranges, thus reducing the amount of data transferred to the destination server.

[0061] Once the VM_MIGRATION_ACTIVE flag is set by the source hypervisor, the SSD controller may log the LBA ranges written to a LBA_DIRTY_LOG. Such a log may be used by the source to find and copy over the delta data changes done by the VM during the brownout phase. That is to say, the source hypervisor may copy all the VM data initially when VM is still active, and then copy only the differences after pausing the VM for migration. Using such a step data copy process, the source hypervisor may ensure the data integrity on the destination server when the VM is relaunched.

[0062] NVMe Get Features—VM MIGRATION ACTIVE command may be used by the source hypervisor to read the current values of the parameters of the VM MIGRATION ACTIVE feature.

The source may get, for example, the following parameter values: VM_MIGRATION_ACTIVE; LBA_WRITTEN_THRESHOLD; and LBA_WRITTEN_THRESHOLD_EXCEEDED.

[0063] NVMe Set Features—VM STORAGE STATE command extension may be used by the destination hypervisor to provide the VM storage state values to the SSD controller on the destination SSD. The SSD controller on the destination SSD may plug these values into the appropriate functional layers such as PCIe transport, NVMe host layer, and FTL. By doing so, the destination SSD controller may become ready or close to ready for storage I/O commands from the migrated VM. The source VM storage state may be obtained by the source hypervisor from the source SSD using a Get Log Page command described elsewhere. The source hypervisor may then communicate these state values to the destination server.

[0064] NVMe Get Log Page—VM STORAGE STATE command may be used by the source hypervisor to read the VM storage state from the attached SSD. The various possible parameters and fields of VM storage state are described elsewhere. After reading the VM storage state log, the source hypervisor may send that log to the destination hypervisor. The destination hypervisor may then use the new Set Features—VM STORAGE STATE command to provide the relevant storage parameters and values to the SSD attached to the destination server.

[0065] NVMe Get Log Page—VM DIRTY STATE command may be used by the source hypervisor to read the delta changes happened in the VM storage state after the source hypervisor sets the VM_MIGRATION_ACTIVE flag. The source hypervisor may use this command after the VM on the source server is paused or frozen. The source hypervisor may initially copy the whole storage state while VM is running, and then copy only the storage state differences after the VM is stopped. The delta storage state may be sent to the destination hypervisor which in turn may use the Set Features—VM STORAGE STATE command to reflect the final VM storage state on the destination server SSD.

[0066] NVMe Get Log Page—VM DIRTY LBA RANGES command may be used by the source hypervisor to read the LBAs written or modified by the VM during the brownout phase. In the brownout phase, the source hypervisor may copy all the LBAs allocated and/or data written by the VM to the destination server while the VM is concurrently running. Since the VM is active during this phase, the VM may write and modify some of the storage LBAs. Hence, during the blackout phase—after the VM is stopped on the source server—the source hypervisor may copy the modified or dirty LBAs to the destination server. The SSD may track and create a dirty LBA ranges log after the hypervisor sets the VM_MIGRATION_ACTIVE flag. The SSD controller may use the LBA_WRITTEN_THRESHOLD to record the LBA ranges updated by the VM. If the number of LBAs written by the VM is greater than the threshold set by the source hypervisor, the SSD may set the LBA_WRITTEN_THRESHOLD_EXCEEDED flag. After reading this log, the source hypervisor may use the LBA ranges to read out that data and send that data to the destination server for persistently storing in the SSD attached to the destination server over its lifetime as well. If the source hypervisor finds that the LBA_WRITTEN_THRESHOLD_EXCEEDED flag is set in this log page, the source hypervisor may choose to ignore the dirty LBA ranges and just copy the whole LBA space.

[0067] The Identify—VM MIGRATION ASSIST SUPPORT data structure may be used by the SSD controller to indicate VM Live Migration Assist support to the source as well as destination hypervisors. The unused fields of the Identify data structure may be used for this purpose.

[0068] All of the above described commands are examples of some embodiments of the inventive concept. Other embodiments of the inventive concept may use different commands or create new custom vendor-defined commands to achieve the same functionality.

Source Hypervisor Operational Flow

[0069] At a high level, live VM migration from the source server to the destination server is done in three phases. After the migration decision, in the first, initial phase, the migration planning and provisioning of resources may be done. In the second phase (aka brownout phase), the source

hypervisor may copy the current state and values of the VM resources such as memory, network, and storage to the destination server. In the final phase (aka blackout) phase, the source hypervisor may stop or pause the VM and copy the delta differences from the earlier copy to the destination sever, after which the VM may be relaunched on the destination server.

[0070] For copying the storage state and storage data from the source server to the destination server, the source hypervisor may perform a series of steps. First during the brownout phase, while the VM is still running, the source hypervisor may set the VM_MIGRATION_ACTIVE flag in the SSD. The source hypervisor may then initiate the first copy of the storage state and stored data. The source hypervisor may read the current storage state using Get Log page command. The read storage state may then be passed to the destination server for restoring it there. After that, the source hypervisor may decide whether to copy certain LBA ranges to the destination or move the whole capacity allocated to that VM. The source hypervisor may use the Get Features command to fetch the LBA ranges written by the VM over its lifespan. If the LBA_WRITTEN_THRESHOLD_EXCEEDED flag is set, hypervisor may read the whole capacity and send that data to the destination server. If the LBA_WRITTEN_THRESHOLD_EXCEEDED flag is not set, then the source hypervisor may selectively read the data from LBA ranges written by the VM and send that data to the destination server. After all the VM resource states and data are moved to the destination server, the source hypervisor may stop the VM. The source hypervisor may then use the Get Log Page commands to read the delta modifications to the storage state and the LBA data. The source hypervisor may then move the differences in state and data to the destination server. At that point the whole state and data values of the VM being migrated may be restored and ready on the destination server. The destination server may relaunch the VM on that server.

Destination Hypervisor Operational Flow

[0071] The destination hypervisor may perform the restoration of the VM resource states on the destination server. For successful VM migration, the VM resource state and data may be recreated on the destination server. By having the VM resource state and data on the destination server identical or very close to the VM resource state and data on the source server, the applications running on the migrated VM may not be impacted. That is to say, the objective of the live VM migration is to make the migration process transparent to the actual applications.

[0072] The destination hypervisor may receive the state of the VM resources on the source server. For storage state, the destination hypervisor may use the Set Features command to provide the VM STORAGE STATE to the attached SSD. The destination hypervisor may receive all the VM data from the source server. The destination hypervisor may persist all the received data from the source server. The destination hypervisor may then receive the delta differences in the VM state and data. The destination hypervisor may then persist the modified data and perform any reconfigurations of the resources based on the delta state changes on the source server. At that point, the destination server may replicate the state of the VM resources and VM data. That is to say, at that point the destination hypervisor may be ready to launch the migrating VM on its server.

[0073] In another embodiment of the inventive concept, the VM storage state and data migration mechanism may also be utilized to take snapshot images and work as an archive, backup and restore mechanism.

[0074] FIG. 1 shows a machine including a storage device equipped to assist in live migration, according to an embodiment of the inventive concept. In FIG. 1, machine 105, which may also be termed a host, is shown. Machine 105 may include processor 110. Processor 110 may be any variety of processor. (Processor 110, along with the other components discussed below, are shown outside machine 105 for ease of illustration: embodiments of the inventive concept may include these components within machine 105.) While FIG. 1 shows a single processor 110 in machine 105, machine 105 may include any number of processors, each of which may be single core or multi-core processors, each of which may implement a Reduced Instruction Set Computer (RISC)

architecture or a Complex Instruction Set Computer (CISC) architecture (among other possibilities), and may be mixed in any desired combination.

[0075] Machine **105** may also include memory **115**. Memory **115** may be any variety of memory, such as flash memory, Dynamic Random Access Memory (DRAM), Static Random Access Memory (SRAM), Persistent Random Access Memory, Ferroelectric Random Access Memory (FRAM), or Non-Volatile Random Access Memory (NVRAM), such as Magnetoresistive Random Access Memory (MRAM) etc. Memory **115** may also be any desired combination of different memory types. Machine **105** may also include memory controller **120**, which may be used to manage access to memory **115**.

[0076] Machine **105** may also include a storage device. Storage device **125** may be used to store data. Processor **110** may run device driver **130**, which may support access to storage device **125**. Embodiments of the inventive concept may include any desired storage device, which may operate using any desired storage principle. Thus, the storage device may be a Solid State Drive (SSD), a hard disk drive, or any other desired storage device, and may store data using sector-based storage, block-based storage, or key-value based storage, among other possibilities. While FIG. **1** shows only one storage device **125**, embodiments of the inventive concept may support any number of installed storage devices in machine **105**.

[0077] Although FIG. **1** depicts machine **105** as a server (which could be either a standalone or a rack server), embodiments of the inventive concept may include machine **105** of any desired type without limitation. For example, machine **105** could be replaced with a desktop or a laptop computer or any other machine that may benefit from embodiments of the inventive concept. Machine **105** may also include specialized portable computing machines, tablet computers, smartphones, and other computing machines. In addition, an application that may be accessing data from storage device **125** may be located in another machine, separate from machine **105** and accessing machine **105** via a network connection traversing one or more networks of any types (wired, wireless, global, etc.).

[0078] FIG. **2** shows additional details of machine **105** of FIG. **1**. In FIG. **2**, typically, machine **105** includes one or more processors **110**, which may include memory controllers **120** and clocks **205**, which may be used to coordinate the operations of the components of device **105**. Processors **110** may also be coupled to memories **115**, which may include random access memory (RAM), read-only memory (ROM), or other state preserving media, as examples. Processors **110** may also be coupled to storage devices **125**, and to network connector **210**, which may be, for example, an Ethernet connector or a wireless connector. Processors **110** may also be connected to buses **215**, to which may be attached user interfaces **220** and Input/Output interface ports that may be managed using Input/Output engines **225**, among other components.

[0079] FIG. **3** shows machine **105** of FIG. **1** and storage device **125** of FIG. **1** supporting virtual machines (VMs). When machine **105** supports VMs, each VM operates as though it was a physically separate machine, and is not aware that other VMs may be running on machine **105**. (This is not to say that a virtual machine is not affected by the operations of other VMs: when one VM is executing commands another VM may be stalled, waiting for resources of machine **105** to become available.) Thus, while FIG. **3** shows machine **105** supporting only VM **305**, embodiments of the inventive concept may include any number-zero or more-VMs on machine **105**.

[0080] Each VM may expect access to various resources of machine **105**, such as processor **110** of FIG. **1**, memory **115** of FIG. **1**, and storage device **125** of FIG. **1**. Thus, VM **305** may be thought of as including application **310**, which may run on system software stack **315**, which may sit on top of operating system **320** (an operating system local to VM **305**, and not necessarily shared across multiple VMs). VM **305** may also include memory **325**, which may be a portion of memory **115** of FIG. **1** (or, in some embodiments of the inventive concept in which at least a portion of memory **115** of FIG. **1** is available to more than one VM, memory **325** may be the entirety of memory **115**). VM **305** may also include network driver **330** to access a network (which may be a local area

network (LAN), a wide area network (WAN), a metropolitan area network (MAN), a global network (such as the Internet), or a combination thereof, and which may be wired, wireless, or a combination thereof), and storage driver **335** to access a storage device such as storage device **125**. [0081] To support VM **305**, machine **105** may include hypervisor **340**. Hypervisor **340** may be thought of as an operating system to manage VMs. Each VM communicates with hypervisor **340**, which in turn may communicate with a network or storage device **125**.

[0082] Storage device **125** may expose one or more functions, such as physical function (PF) **345**. The exposed functions may be PFs, virtual functions (VFs), or a combination thereof. In addition, as described in U.S. patent application Ser. No. 16,846,271, filed Apr. 10, 2020, now pending, which claims the benefit of U.S. Provisional Patent Application Ser. Nos. 62,865,962, filed Jun. 24, 2019 and U.S. Provisional Patent Application Ser. No. 62,964,114, filed Jan. 21, 2020, all of which are incorporated by reference herein for all purposes, storage device **125** may be associated with a lightweight bridge, which may present PFs to machine **105** on behalf of storage device **125**. Thus, the PFs and/or VFs may be exposed for storage device **125** by storage device **125** itself or by another component on behalf of storage device **125**.

[0083] For each function, such as PF **345**, storage device may have an associated Non-Volatile Memory Express (NVMe) controller, such as NVMe controller **350**. NVMe controller **350** may manage operations within storage device **125** for requests sent via function **345**. In some embodiments of the inventive concept, each exposed function may have its own NVMe controller; in other embodiments of the inventive concept, a single NVMe controller may handle requests on behalf of two or more exposed functions.

[0084] In some embodiments of the inventive concept, each VM uses a unique function for its requests of storage device **125**. Put another way, when a VM, such as VM **305**, sends a request intended for storage device **125**, hypervisor **340** may receive the request and direct that request to the function associated with that VM. (Thus, hypervisor **340** may store a mapping between individual functions exposed by storage device **125** and VMs running on machine **105**, to facilitate communication passing correctly between VMs and storage device **125**.) This option provides a simple way for storage device **125** to know which VM is issuing the request, since each VM may use a different function. So, for example, if a storage device offers **16** exposed functions, 16 different VMs may each use one of the exposed functions. But if each VM uses a unique exposed function, then the total number of VMs that may be supported by the machine **105** may be capped at the total number of functions exposed by all installed storage devices. Thus, for example, if machine **105** has a total of three storage devices installed, two of which expose eight functions and the third exposing **16** functions, the total of VMs that may be supported by machine **105** may be capped at **32** (eight plus eight plus sixteen).

[0085] In other embodiments of the inventive concept, multiple VMs may share a function for requests of storage device **125**. Put another way, multiple VMs may send requests to hypervisor **340**, and some (or all) of the requests from these VMs may be sent to a shared function exposed by storage device **125**. Embodiments of the inventive concept may include all functions exposed by storage device **125** being shared by at least two VMs, each function exposed by storage device **125** used by a single VM, or any desired combination of the two approaches.

[0086] Note that in embodiments of the inventive concept where a single function, such as function **345**, is shared by two or more VMs, storage device **125** might not automatically have enough information to know what data is actually stored for a given VM. That is, for any given write request sent to storage device **125** via function **345**, storage device **125** may not know what VM issued that write request. Thus, for example, if storage device **125** has stored, say, approximately 1 GB of data received via function **345**, storage device **125** may not know how much of that data was written by a VM now being migrated (as compared with data written by a VM not being migrated).

[0087] There are ways to inform storage device **125** what VM issued a particular write request received across a function shared by multiple VMs. For example, hypervisor **340** may add a tag (as

metadata) that identifies which VM issued the write request. (Including such a tag may be useful to hypervisor **340** as well, to support returning data read from storage device **125** to the appropriate VM, rather than having hypervisor **340** identify the read request for which the data was sent by storage device **125**, say, by examining the LBA (or some other metadata) associated with the data.) In that manner, storage device **125** may use the tag to track what data belongs to what VM, even when two or more VMs share an exposed function of storage device **125**.

[0088] One example of such a tag that may be used to identify data is to assign write requests to streams. For example, hypervisor **340** may assign write requests from different VMs to different streams. Then, storage device **125** may use the stream identifiers to track data written by different VMs. Streaming of data may be used, for example, by storage device **125** (when storage device **125** is an SSD) to write all data associated with a particular stream ID to a set of blocks that contains no data from other streams. Streaming of data may benefit garbage collection in SSDs: if a VM is to be migrated, then all data associated with that VM may be moved from the SSD to another storage device, and all the blocks occupied by that data may be invalidated and subject to garbage collection without leaving any data in those blocks that require programming to other blocks in the SSD. Note that the term “stream ID” as used herein may identify a stream within the SSD, or it may identify a stream within machine **105** (and the SSD may be left to manage mapping the management of stream IDs within machine **105** to streams within the SSD).

[0089] Examples of how streaming may be managed may be found in: U.S. patent application Ser. No. 15/089,237, filed Apr. 1, 2016, now U.S. Pat. No. 9,959,046, issued May 1, 2018, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/273,323, filed Dec. 30, 2015; U.S. patent application Ser. No. 15/146,708, filed May 4, 2016, now U.S. Pat. No. 9,898,202, issued Feb. 20, 2018, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/302,162, filed Mar. 1, 2016 and which is a continuation-in-part of U.S. patent application Ser. No. 15/046,439, filed Feb. 17, 2016, now U.S. Pat. No. 9,880,780, issued Jan. 30, 2018, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/261,303, filed Nov. 30, 2015; U.S. patent application Ser. No. 15/690,270, filed Aug. 29, 2017, currently pending, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/501,087, filed May 3, 2017, and U.S. Provisional Patent Application Ser. No. 62/511,957, filed May 26, 2017; U.S. patent application Ser. No. 15/821,708, filed Nov. 22, 2017, now allowed, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/561,674, filed Sep. 21, 2017 and which is a continuation-in-part of U.S. patent application Ser. No. 15/620,814, filed Jun. 12, 2017, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/490,027, filed Apr. 25, 2017; U.S. patent application Ser. No. 16/775,262, filed Jan. 28, 2020, currently pending, which is a divisional of U.S. patent application Ser. No. 15/167,974, filed May 27, 2016, now U.S. Pat. No. 10,592,171, issued Mar. 17, 2020, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/309,446, filed Mar. 16, 2016; U.S. patent application Ser. No. 16/856,020, filed Apr. 22, 2020, currently pending, which is a continuation of U.S. patent application Ser. No. 15/499,877, filed Apr. 27, 2017, now U.S. Pat. No. 10,656,838, issued May 19, 2020, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/458,566, filed Feb. 13, 2017, and U.S. Provisional Patent Application Ser. No. 62/471,350, filed Mar. 14, 2017 (U.S. patent application Ser. No. 15/499,877, filed Apr. 27, 2017 is also a continuation-in-part of U.S. patent application Ser. No. 15/344,422, filed Nov. 4, 2016, now U.S. Pat. No. 10,282,324, issued May 7, 2019, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/383,302, filed Sep. 2, 2016, and which is a continuation-in-part of U.S. patent application Ser. No. 15/144,588, filed May 2, 2016, currently pending, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/245,100, filed Oct. 22, 2015 and U.S. Provisional Patent Application Ser. No. 62/192,045, filed Jul. 13, 2015; U.S. patent application Ser. No. 15/499,877, filed Apr. 27, 2017 is also a continuation-in-part of U.S. patent application Ser. No. 15/090,799, filed Apr. 5, 2016, now U.S. Pat. No. 10,509,770, issued Dec. 17, 2019, which claims the benefit of U.S. Provisional Patent

Application Ser. No. 62/245,100, filed Oct. 22, 2015 and U.S. Provisional Patent Application Ser. No. 62/192,045, filed Jul. 13, 2015); U.S. patent application Ser. No. 16/866,545, filed May 4, 2020, currently pending, which is a continuation of U.S. patent application Ser. No. 16/219,936, filed Dec. 13, 2018, now allowed, which is a continuation of U.S. patent application Ser. No. 15/458,968, filed Mar. 14, 2017, now U.S. Pat. No. 10,216,417, issued Feb. 26, 2019, which claims the benefit of U.S. Provisional Patent Application Ser. No. 62/448,958, filed Jan. 20, 2017, and U.S. Provisional Patent Application Ser. No. 62/413,177, filed Oct. 26, 2016, all of which are incorporated by reference herein.

[0090] NVMe controller **350** may communicate with host interface **355**. Each function **345** may have its own NVMe controller **350**. Host interface **355** may act as a general interface for requests received from machine **105**. Host interface **355** may communicate with flash translation layer (FTL) **360**, which may handle the translation of logical identifiers (or logical addresses) received from machine **105** into physical addresses on storage device **125** where the data is actually stored. These logical identifiers do not need to actually be addresses: for example, while these logical identifiers may be logical block addresses (LBAs) for a block-based storage device, they may be keys for a key-value storage device. This information may be passed to flash channels **365**, which may be used to access the data from flash chips **370**. (The use of flash channels **365** and flash chips **370** suggest that storage device **125** in FIG. **3** may be an SSD, flash channels **365** and flash chips **370** may be replaced with other comparable components for storage devices of other types, such as hard disk drives.)

[0091] While FIG. **3** shows an abstracted view of storage device **125** as part of machine **105** supporting VMs, FIG. **4** shows another view of storage device **125** of FIG. **1**. In FIG. **4**, SSD **125** may include host interface **355** (which may also be termed a “host interface logic”, or “HIL”), SSD controller **405**, and various flash memory chips **370-1** through **370-8** (also termed “flash memory storage”), which may be organized into various channels **365-1** through **365-4**. Host interface logic **355** may manage communications between SSD **125** and other components (such as processor **110** of FIG. **1**). In some embodiments of the inventive concept, SSD **125** may also support its own network connection (as opposed to passing through a network interface that is part of machine **105** of FIG. **1**): in such embodiments of the inventive concept, host interface logic **355** may also manage communications with devices remote from SSD **125**: that is, devices that are not considered part of machine **105** of FIG. **1**, but in communication with SSD **125**. These communications may include read requests to read data from SSD **125**, write requests to write data to SSD **125**, and delete requests to delete data from SSD **125**. Host interface logic **355** may manage an interface across only a single port, or it may manage interfaces across multiple ports.

Alternatively, SSD **125** may include multiple ports, each of which may have a separate host interface logic **355** to manage interfaces across that port. Embodiments of the inventive concept may also mix the possibilities (for example, an SSD with three ports might have one host interface logic to manage one port and a second host interface logic to manage the other two ports). To support direct communication with remote devices over a network, HIL **355** may include an Ethernet component (not shown in FIG. **4**) or some equivalent network communication component.

[0092] SSD controller **405** may manage the read and write operations, along with garbage collection and other operations, on flash memory chips **370-1** through **370-8** using a flash memory controller (not shown in FIG. **4**). SSD controller **405** may include flash translation layer **360**, storage state **410**, and VM Migration State Monitor and Capture Module **415**. Flash translation layer **420**, as discussed above with reference to FIG. **3**, may manage translation between logical data identifiers as used by machine **105** of FIG. **1** and the physical locations in flash chips **370-1** through **370-8** where the data represented by the logical data identifiers is actually stored.

[0093] Both NVMe controller **350** of FIG. **3** and SSD controller **405** are described as being used in embodiments of the inventive concept: there may be confusion about the presence of both controllers. SSD controller **405** may manage reading data from and writing data to flash chips **370-**

1 through **370-8**; NVMe controller **350** of FIG. **3**, on the other hand, may act as part of the interface for the VM. Their operations are related, but not entirely overlapping.

[0094] While SSD controller **405** is typically implemented using hardware in some form (for example, as a custom controller for SSD **125** such as an Application-Specific Integrated Circuit (ASIC), although SSD controller **405** may be implemented using a general purpose processor of some sort (such as a central processing unit, a graphics processing unit (GPU), or a general purpose GPU (GPGPU)) with appropriate software or using a suitably configured Field Programmable Gate Array (FPGA), among other possibilities), the implementation of NVMe controller **350** of FIG. **3** may be more varied. In some embodiments of the inventive concept, NVMe controller **350** of FIG. **3** may be implemented using hardware similar to (but separate from) SSD controller **405**. Thus, NVMe controller **350** of FIG. **3** may be implemented using hardware, such as an ASIC or a suitably configured FPGA. In other embodiments of the inventive concept, NVMe controller **350** of FIG. **3** may be implemented using software (that may run on available appropriate hardware). In yet other embodiments of the inventive concept, NVMe controller **350** of FIG. **3** may include both hardware and software elements (such as a general purpose processor of some sort running suitable software).

[0095] In some embodiments of the inventive concept, NVMe controller **350** of FIG. **3** may be implemented separately from (but communicatively coupled to) SSD controller **405** (perhaps as part of host interface **355**, perhaps separately from host interface **355** as well); in other embodiments of the inventive concept, NVMe controller **350** of FIG. **3** may be implemented as part of SSD controller **405**. NVMe controller **350** of FIG. **3** may also be implemented externally to storage device **135**. (The idea that NVMe controller **350** of FIG. **3** may be implemented externally to storage device **125** may seem contrary to the description of NVMe controller **350** of FIG. **3** supporting a particular function exposed by storage device **125**. But U.S. patent application Ser. No. 16,846,271, filed Apr. 10, 2020, now pending, which claims the benefit of U.S. Provisional Patent Application Ser. Nos. 62,865,962, filed Jun. 24, 2019 and U.S. Provisional Patent Application Ser. No. 62,964,114, filed Jan. 21, 2020, all of which are incorporated by reference herein for all purposes, describes how a function may be exposed on behalf of storage device **125** using a lightweight bridge that may be external to the storage device itself. NVMe controller **350** of FIG. **3** may therefore be implemented as part of this lightweight bridge or as part of a component that may be present somewhere one of the exposed functions of the lightweight bridge and storage device **125**.) In embodiments of the inventive concept where NVMe controller **350** of FIG. **3** may be external to storage device **125**, NVMe controller **350** of FIG. **3** may communicate with multiple different storage devices **125** (for example, data for a single VM may span multiple storage devices, or data for one VM may be written to one storage device while data for another VM may be written to another storage device). In such embodiments of the inventive concept, NVMe controller **350** of FIG. **3** may store information associating which storage device may store a particular piece of data (similar to the flash translation table, but mapping LBAs or other logical identifiers to a storage device rather than an address within the storage device).

[0096] Storage state **410** include information about the state of SSD **125**. As discussed further below with reference to FIG. **6**, storage state **410** may include information about the configuration of the transport protocol(s) used for communications between machine **105** of FIG. **1** and SSD **125**, information about the configuration of the interface (that is, PF(s) **345** of FIG. **3**) across which communications between machine **105** of FIG. **1** and SSD **125** travel, information about Input/Output (I/O) activity associated with individual VMs, and information about the flash translation layer for individual VMs. Note that storage state **410** may represent the state of the entirety of SSD **125**, or it may represent the state of SSD **125** as associated with a particular VM (since different VMs may have different configurations), in which case there may be multiple storage states **410** stored in SSD **125** (or a single storage state **410**, but with information delineated by VM). Storage state **410** may be stored in any desired form of storage, such as some form of

memory (for example, RAM or equivalents) or non-volatile storage, and may be the same storage as that used to store the data (for example, flash chips **370-1** through **370-8**) or a different medium. [0097] VM Migration State Monitor and Capture Module **415** may manage capturing information about storage state **410** for an individual VM that is to be migrated from machine **105** of FIG. **1** to another VM. VM Migration State Monitor and Capture Module **415** may also assist in the transfer of data to the destination storage device. In addition, VM Migration State Monitor and Capture Module **415** may assist with configuring SSD **125** for a VM that is being migrated to machine **105** of FIG. **1**. When assisting with receiving a migrating VM, VM Migration State Monitor and Capture Module **415** may receive state information from the source storage device (directly from the source storage device or indirectly from hypervisor **340** on either the source or destination machine **105** of FIG. **1**) and configure SSD **125** accordingly. VM Migration State Monitor and Capture Module **415** may interface with host interface logic **355**, flash translation table **360**, cache **420** (as well as cache policy **425** and/or prefetch policy **425**), and flash channels **365-1** through **365-4**.

[0098] SSD controller **405** may also include cache **420** and cache policy **425** and/or prefetch policy **425**. Cache **420** may be used to store data in faster storage than flash chips **370-1** through **370-8**: for example, some form of memory such as RAM or equivalents. For example, cache **420** may be used to buffer data to be written to flash chips **370-1** through **370-8**, or may store data written recently by machine **105** of FIG. **1**, in the expectation that such data may be needed again soon. Alternatively, cache **420** may be used to prefetch data from flash chips **370-1** through **370-8** that SSD **125** expects will be requested soon. Note that the two uses of cache **420** are compatible: cache **420** may support both a cache policy for data received from machine **105** of FIG. **1** and a prefetch policy for data expected to be requested by machine **105** of FIG. **1**. Further, in some embodiments of the inventive concept, the VM may have some control over the cache policy and/or prefetch policy, in which case SSD **125** may include multiple cache policies and/or prefetch policies as requested by different VMs storing data on SSD **125**.

[0099] When a VM is being migrated, cache policy **425** and/or prefetch policy **425** may be modified to better support VM migration. For example, data received from machine **105** of FIG. **1** may be kept in cache **420** longer than normal according to the cache policy, since SSD **125** may expect the data to be requested shortly to be sent to the destination storage device. Similarly, the prefetch policy may be modified to request more data stored for the VM, again with the expectation that the data will be requested shortly to be sent to the destination storage device. The prefetching of VM data may also trigger certain data processing functions such as error correction and decryption (as needed), so that such data processing functions may be completed or at least started earlier in the expectation of that data read request.

[0100] While FIG. **4** shows SSD controller **405** as including flash translation layer **360**, storage state **410**, VM Migration State Monitor and Capture Module **415**, cache **420**, and cache policy **425** and/or prefetch policy **425**, embodiments of the inventive concept may locate these modules in any desired locations. Embodiments of the inventive concept may also locate these modules within different portions of SSD **125**: for example, none of these modules might be within SSD controller **405**.

[0101] While FIG. **4** shows SSD **125** as including eight flash memory chips **370-1** through **370-8** organized into four channels **365-1** through **365-4**, embodiments of the inventive concept may support any number of flash memory chips organized into any number of channels. Similarly, while FIG. **4** shows the structure of a SSD, other storage devices (for example, hard disk drives) may be implemented using a different structure, but with similar potential benefits.

[0102] FIG. **5** shows a virtual machine being migrated from one machine to another. In FIG. **5**, VM **305** is operating on machine **105-1**, and is to be migrated to machine **105-2**. VM **305** has associated data (including state information) stored on source SSD **125-1**, which may be migrated with VM **305**, intended for destination SSD **125-2**. Once this migration is complete, VM **305** will be

operating on machine **105-2**, ideally with no evidence that VM **305** was previously operating on machine **105-1**.

[0103] Up until now, VM migration has been discussed in the abstract, without discussing the specifics of how VM migration is performed, or how a particular VM might be selected for migration. While the process by which a particular VM may be selected for migration is beyond the scope of this document, there are numerous reasons why a VM might be migrated. As data centers establish new servers, VMs may be migrated to those new servers to utilize their resources (and to lessen the burden on the old server). A particular server might require maintenance that would take it offline, in which case VMs running on that server may be migrated so that they may continue to execute. A particular server may be old and need to be retired, in which case VMs executing on that server should be migrated to other servers. The loads on different servers might vary, so that one server is heavily loaded and another is lightly loaded: load balancing might justify migrating a VM from the heavily loaded server to the lightly loaded server. A VM might change its requirements, in which case the server on which the VM is executing might be insufficient to the new requirements and a new server should be used. And there might be other reasons why a VM might be selected for migration.

[0104] VM migration itself is a complicated process, and may require significant processing and resource utilization. Performing a VM migration may therefore impact the performance of the VM itself and other VMs installed on the machine. Thus, it is desirable that VM migration be performed in a manner that minimizes this load, so that user applications are minimally affected.

[0105] VM migration should also be transparent to the VM itself. That is, the VM should not be aware of which server it is installed in, or the fact that the VM is being (or has been) migrated. A performance delay while the VM is being migrated is therefore undesirable.

[0106] The process of migrating a VM is typically performed in three phases. In the first phase, called the preparation phase, destination machine **105-2** is prepared to receive VM **305**. In the preparation phase, resources may be allocated on destination machine **105-2**. These resources may include an assigned core in processor **110** of FIG. **1**, a section of memory **115** of FIG. **1** for VM **305**, a portion of storage space on source storage device **125-1**, an exposed function **345** of FIG. **3** to be used by VM **305**, and so on. Note that by the time the preparation phase begins, VM **305** has been selected for migration, as has destination machine **105-2** to which VM **305** is to be migrated.

[0107] In the second phase, called the brownout phase, hypervisor **340-1** begins to migrate storage state **410** of FIG. **5** and the data for VM **305** to destination machine **105-2**. During the brownout phase, VM **305** is still executing and may access its data and its resources. Thus, during the brownout phase, some information relating to VM **305** may change. These changes may include data as stored on source storage device **125-1**, changes to the configuration of source storage device **125-1**, data in memory **325** of FIG. **3**, and so on. Thus, while most of VM **305** may be transferred to destination machine **105-2** during the brownout phase, it is reasonable to expect that some data will have changed but not yet migrated.

[0108] To address the changes that were made to VM **305** during the brownout phase, VM migration includes a third phase, called the blackout phase. In the blackout phase, VM **305** is halted (other terms that might be used to describe what happens to VM **305** during the blackout phase include suspension, pausing, stoppage, freezing, etc.). By halting VM **305**, VM **305** is prevented from making any further changes to its data. During the blackout phase, hypervisor **340-1** may migrate any changes or deltas made to VM **305** relative to the copy made during the brownout phase. Once all changes or deltas have been copied to destination machine **105-2**, VM **305** as stored on both source machine **105-1** and destination machine **105-2** should be identical. VM **305** may then be started on destination machine **105-2**, and the copy of VM **305** on source machine **105-1** may be terminated, releasing the resources used by VM **305** on source machine **105-1**.

[0109] As noted above, VM migration is expensive in terms of processing and resources. Since hypervisor **340-1** handles most of VM migration (there are corresponding operations being

performed by hypervisor **340-2** on destination machine **105-2**), hypervisor **340-1** has fewer resources available to be used by other VMs not being migrated. By shifting some of the work regarding VM migration to source storage device **125-1**, the load on hypervisor **340-1** may be reduced, to the benefit of all VMs on source machine **105-1**.

[0110] Since source storage device **125-1** does not have access to everything associated with VM **305** (for example, information about applications **310** of FIG. 3 running on VM **305**, or the state of memory **325** of FIG. 3), source storage device **125-1** may not completely take over VM migration. Nor is VM migration necessarily limited to just transferring the data stored on source storage device **125-1**. But with storage devices capable of storing terabytes of data, source storage device **125-1** may manage the lion's share of data about a VM. Having source storage device **125-1** involved in migrating information from source storage device **125-1** may significantly reduce the load on hypervisor **340-1**. Further, since source storage device **125-1** may track what requests are received over each exposed function, source storage device **125-1** may readily determine what data is stored for a VM being migrated. Hypervisor **340-1** may not necessarily know what data stored on source storage device **125-1** is for VM **305** (which is being migrated), and may have to request such information from source storage device **125-1** (via function **345** of FIG. 3). Similarly, source storage device **125-1** may have ready access to information such as the number of logical identifiers used by VM **305**, or information about the translation table for those logical identifiers, that may be hidden from or unknown by hypervisor **340-1**.

[0111] Thus, at a request from hypervisor **340-1**, source storage device **125-1** may begin collecting information that will assist hypervisor **340-1** in performing VM migration. This request may set a flag within source storage device **125-1**. Based on exposed function **345** of FIG. 3, source storage device **125-1** may know which VM is being migrated. Source storage device **125-1** may then start to collect information that may be used in the VM migration. For example, source storage device **125-1** may collect storage state **410** of VM **305**, and source storage device **125-1** may begin identifying all logical identifiers (such as LBAs or keys) used by VM **305** to write data. Source storage device **125-1** may also modify the cache policy and/or prefetch policy to benefit the VM migration. That is, source storage device **125-1** may modify the cache policy to retain data in the cache **420** of FIG. 4 rather than flushing the data to the storage medium, and/or source storage device **125-1** may modify the prefetch policy to favor reading data written for VM **305** into cache **420** of FIG. 4. Either or both of these policy modifications may expedite the process of transferring data for VM **305** to destination storage device **125-2**.

[0112] At the request of hypervisor **340-1**, source storage device **125-1** may generate a log page or other report that includes storage state **410** of FIG. 4. Source storage device **125-1** may send this log page or other report to hypervisor **340-1**, which may then send that information to hypervisor **340-2** on destination machine **105-2**, which in turn may use storage state **410** of FIG. 4 to configure destination storage device **125-2**. In this manner, destination storage device **125-2** may ultimately be configured to mirror the configuration of source storage device **125-1**.

[0113] Alternatively, source storage device **125-1** may send storage state **410** of FIG. 4 (again, as a log page or other report) directly to destination storage device **125-2**. Destination storage device **125-2** may then configure itself according to storage state **410** of FIG. 4. For source storage device **125-1** to send storage state **410** of FIG. 4 directly to destination storage device **125-2**, source storage device **125-1** may use its own network connection (such as an Ethernet connection) or a network connection of source machine **105-1**.

[0114] Source storage device **125-1** may also determine whether a threshold number of LBAs (or other logical identifiers) have been written to storage device. This determination may expedite the process of transferring the data from source storage device **125-1** to destination storage device **125-2**.

[0115] To explain how this determination may be helpful, consider source storage device **125-1** with approximately 100 gigabytes of data storage capacity allocated for VM **305**. If VM **305** has

written, say, only approximately four gigabytes of data to source storage device **125-1**, then it may be more efficient for hypervisor **340-1** to identify those LBAs that have been written and transfer them to directly destination storage device **125-2** (say, by requesting those LBAs from source storage device **125-1**, then sending read requests for those LBAs to source storage device **125-1**), or by sending them to hypervisor **340-2** on destination machine **105-2**. But if, say, approximately 98 gigabytes of data have been written to source storage device **125-1** for VM **305**, it might be more efficient to simply transfer the entirety of the allocated space for VM **305** from source storage device **125-1** to destination storage device **125-2**.

[0116] To put the entire process into context, in the preparation phase of VM migration, hypervisor **340-1** of source machine **105-1** may first send a request to source storage device **125-1**, letting source storage device **125-1** know that VM **305** is to be migrated to destination machine **105-2**. Hypervisor **340-1** of source machine **105-1** may use this request to set a flag, which may be a labeled VM_MIGRATION_ACTIVE flag. Once this flag is set, source storage device **125-1** may start preparing for VM **305** to be migrated from source machine **105-1** to destination machine **105-2**. This preparation may include, for example, identifying storage state **410** of FIG. 4 and/or modifying the cache policy and/or prefetch policy to favor data for VM **305** in cache **420** of FIG. 4. Hypervisor **340-2** of destination machine **105-2** may also inform destination storage device **125-2** that a VM migration is underway, so that destination storage device **125-2** may be configured for VM **305** and to receive the data for VM **305**.

[0117] Eventually, hypervisor **340-1** of source machine **105-1** may send a request for storage state **410** of FIG. 4 to source storage device **125-1**. This request may specify that storage state **410** of FIG. 4 be prepared in a log page or some other report. This request may also specify that storage state **410** of FIG. 4 be made available to hypervisor **340-1** of source machine **105-1** (either by being sent in a message back to hypervisor **340-1** of source machine **105-1** or by storing storage state **410** of FIG. 4 in a location accessible to hypervisor **340-1** of source machine **105-1**), or be sent to destination storage device **125-2** (perhaps over a network). This request may include a GET_LOG_PAGE request for the VM_STORAGE_STATE data. Destination storage device **125-2** may then configure itself according to storage state **410** of FIG. 4 (whether destination storage device **125-2** does this configuration on its own or at the direction of hypervisor **340-2** of destination machine **105-2** may depend on the specific implementation). This request may include a SET_LOG_PAGE request for the VM_STORAGE_STATE data.

[0118] Eventually, hypervisor **340-1** of source machine **105-1** may enter the brownout phase of VM migration. Source storage device **125-1** may compare the number of logical identifiers used for VM **305** and compare that number to a threshold, which may be labeled LBA_WRITTEN_THRESHOLD. Hypervisor **340-1** of source machine **105-1** may set the value for LBA_WRITTEN_THRESHOLD. Based on this comparison, source storage device **125-1** may set a flag, which may be labeled LBA_WRITTEN_THRESHOLD_EXCEEDED, which may indicate whether the number of logical identifiers is large enough to justify copying the entirety of the data for VM **305**, or if the number of logical identifiers is small enough that it may be more efficient to copy just the logical identifiers that have actually be written.

[0119] Hypervisor **340-1** of source machine **105-1** may access the LBA_WRITTEN_THRESHOLD_EXCEEDED flag to determine the appropriate approach to copying data from source storage device **125-1** to destination storage device **125-2**. Hypervisor **340-1** of source machine **105-1** may then read data for VM **305** (either individual logical identifiers or simply request all data for VM **305**) from source storage device **125-1** and send that data to destination hypervisor **340-2** of destination machine **105-2**, to be written to destination storage device **125-2**. Alternatively, source storage device **125-1** may send the data directly to destination storage device **125-2**, which may be performed in any desired manner: for example, source storage device **125-1** may send the data for VM **305** to destination storage device **125-2** as a stream. Destination storage device **125-2** may then store the data for VM **305** (whether destination storage

device **125-2** stores the data for VM **305** on its own or at the direction of hypervisor **340-2** of destination machine **105-2** may depend on the specific implementation).

[0120] Rather than reading data from source storage device **125-1**, hypervisor **340-1** of source machine **105-1** may opt for destination storage device **125-2** to use a snapshot of data for VM **305** on source storage device **125-1**. That is, at some point (either before or after hypervisor **340-1** of source machine **105-1** starts the migration of VM **305** from source machine **105-1** to destination machine **105-2**, and potentially even before VM **305** is selected for migration), a snapshot of data on source storage device **125-1** may be taken. This snapshot may include all data stored on source storage device **125-1**, or just a snapshot of data for VM **305** on source storage device **125-1**. This snapshot may be generated by VM Migration State Monitor and Capture Module **415** of FIG. 4 in the same way that VM Migration State Monitor and Capture Module **415** of FIG. 4 may be used to select the data for VM **305** for migration from source storage device **125-1** to destination storage device **125-2**, except that the data for VM **305** may be stored somewhere rather than transmitted to destination storage device **125-2** to support the migration of VM **305** from source machine **105-1** to destination machine **105-2**.

[0121] During the brownout phase, source storage device **125-1** may track any changes to source storage device **125-1** being made by VM **305**. These changes may include data being written to source storage device **125-1**, data being deleted from source storage device **125-1**, and changes to the configuration of source storage device **125-1**. These changes, or deltas, represent variances from the information provided to destination storage device **125-2** during either the preparation phase or the brownout phase.

[0122] During the blackout phase, hypervisor **340-1** of source machine **105-1** may pause VM **305**. By pausing VM **305**, VM **305** is prevented from making any further changes to its configuration or data (but neither may VM **305** execute on behalf of a user).

[0123] Hypervisor **340-1** of source machine **105-1** may then request the deltas to storage state **410** of FIG. 4. This request may also specify that the deltas to storage state **410** of FIG. 4 be made available to hypervisor **340-1** of source machine **105-1** (either by being sent in a message back to hypervisor **340-1** of source machine **105-1** or by storing storage state **410** of FIG. 4 in a location accessible to hypervisor **340-1** of source machine **105-1**), or be sent to destination storage device **125-2** (perhaps over a network). This request may include a GET_LOG_PAGE for the VM_DIRTY_STATE data. Destination storage device **125-2** may then configure itself according to the deltas to storage state **410** of FIG. 4 (whether destination storage device **125-2** does this configuration on its own or at the direction of hypervisor **340-2** of destination machine **105-2** may depend on the specific implementation). This request may include a SET_LOG_PAGE request for the VM_STORAGE_STATE data.

[0124] Hypervisor **340-1** of source machine **105-1** may then request the deltas to the data on source storage device **125-1** for VM **305**. This request may also specify that the deltas to the data on source storage device **125-1** for VM **305** be made available to hypervisor **340-1** of source machine **105-1** (either by being sent in a message back to hypervisor **340-1** of source machine **105-1** or by storing storage state **410** of FIG. 4 in a location accessible to hypervisor **340-1** of source machine **105-1**), or be sent to destination storage device **125-2** (perhaps over a network). This request may include a GET_LOG_PAGE for the VM_DIRTY_LBA_RANGES data. Destination storage device **125-2** may then store the deltas to the data for VM **305** (whether destination storage device **125-2** stores the deltas to the data for VM **305** on its own or at the direction of hypervisor **340-2** of destination machine **105-2** may depend on the specific implementation).

[0125] Once all the deltas to storage state **410** of FIG. 4 and the data on source storage device **125-1** for VM **305** have been transferred to destination storage device **125-2**, destination storage device **125-2** is ready for VM **305** to be activated on destination machine **105-2**. Hypervisor **340-2** on destination machine **105-2** may then inform destination storage device **125-2** that VM migration is complete, and VM **305** may be started on destination machine **105-2**. Hypervisor **340-1** of source

machine **105-1** may also inform source storage device **125-1** that VM migration is complete, and VM **305** may be terminated on source machine **105-1** (and all information relating to VM **305** may be deleted from source storage device **125-1**), at which point VM migration is complete.

[0126] Storage devices **125-1** and/or **125-2** may store data in a compressed, encoded, and/or encrypted format. For clarity, compression is intended to mean that the data has been subject to an algorithm that typically reduces the amount of space required to store the data, encoding is intended to mean that the data is stored in a different representation than it was provided to storage devices **125-1** and/or **125-2**, and encryption is intended to mean that the data has been subject to an algorithm that typically renders the original content unrecognizable and unrecoverable without an encryption password or encryption key (the term “encryption key” is intended to cover both possibilities, along with any other terms that may represent similar concepts whether or not listed herein). Thus, for example, data may be compressed using an algorithm such as dictionary compression, Prefix Encoding, Run Length Encoding (RLE), Cluster Encoding, Sparse Encoding, and Indirect Encoding, Bit Packing, or a Lempel-Ziv (LZ) compression method; data may be encoded in formats such as binary, hex, decimal, base64, Distinguished Encoding Rules (DER), or Basic Encoding Rules (BER); and data may be encrypted using an algorithm such as private key cryptosystems (for example, the Data Encryption Standard (DES), Triple DES (3DES), Advanced Encryption Standard (AES)) or public key cryptosystems (for example, the Diffie-Hellman key exchange, Rivest-Shamir-Adleman (RSA), and Pretty Good Privacy (PGP)). None of the above examples are intended to be limiting, and embodiments of the inventive concept may use other compression algorithms, encoding schemes, and encryption algorithms as well.

[0127] As an example of encoding, consider binary coded data (BCD). BCD encodes each digit 0-9 in a number separately, which may require 4 bits of data. Thus, the number 1024, which is made up of the four digits 1, 0, 2, and 4, may be represented in BCD as 0001 0000 0010 0100. On the other hand, if treated as a number, 1024 may be represented in binary as 00000100 00000000 (assuming 16-bit representation of integers). Both BCD and binary are valid ways to represent and store the number 1024: they are just different encodings.

[0128] Compression, encoding, and encryption are theoretically orthogonal operations to each other: applying one operation does not prevent either of the other operations from being performed on data as well. (In practice, combining compression and encoding does not provide much benefit over either alone, and encryption may change the data sufficiently that compression and/or encoding may not be of significant benefit. But this fact is not relevant to the operation of embodiments of the inventive concept.) Thus, data stored on source storage device **125-1** may be compressed, encoded, and/or encrypted. Source storage device **125-1** may apply compression, encoding, and/or encryption in any desired order, and may use these approaches in any desired combination. For example, one storage device might compress, then encode, then encrypt data, while another storage device might encrypt, then compress data, but not bother to encode the data.

[0129] Storage devices **125-1** and/or **125-2** may also offer error correction services. For example, storage devices **125-1** and/or **125-2** may use error correcting codes, such as a single error correction double error detection (SECDED) code, to help protect against the possibility of a data error in a read or write request. Storage devices **125-1** and/or **125-2** may use error correcting codes whether or not the data is compressed, encrypted, and/or encoded.

[0130] In some embodiments of the inventive concept, source storage device **125-1** may provide the data for VM **305** exactly as stored. Thus, if the data is stored compressed, encoded, and/or encrypted, the data for VM **305** ultimately received by destination storage device **125-2** may be compressed, encoded, and/or encrypted. Similarly, if the data includes error correcting codes, the data for VM **305** ultimately received by destination storage device **125-2** may include the error correcting codes. If destination storage device **125-2** uses the same compression, encoding, encryption, and/or error correcting code schemes, this may be acceptable (although for encrypted data, source storage device **125-1** may also need to provide the encryption key for encrypted data

for destination storage device **125-2** to be able to ultimately decrypt the data for VM **305** to use the data, if the encryption algorithm was applied by source storage device **125-1** and not VM **305**). But if source storage device **125-1** and destination storage device **125-2** use different compression, encoding, and/or encryption algorithms or a different error correcting code scheme, the data for VM **305** may be processed to remove any processing that destination storage device **125-2** may be unequipped to handle.

[0131] In other embodiments of the inventive concept, source storage device **125-1** may undo any compression, encoding, encryption, and/or error correcting codes of the data for VM **305** before the data for VM **305** is sent for ultimate delivery to destination storage device **125-2**. Thus, source storage device **125-1** may decompress, decode, decrypt, and/or remove error correcting codes from the data for VM **305** before sending the data for VM **305** as part of migrating VM **305** from source machine **105-1** to destination machine **105-2**. Note that the order in which these processes are performed may be pertinent: if data is encrypted first and then compressed, attempting to decrypt the data before decompressing the data is likely to result in the eventual data being completely unusable by VM **305** after migration to destination machine **105-2**. In addition, various embodiments of the inventive concept may undo some, but not all, of these approaches. For example, source storage device **105-1** might first compress and then encrypt the data for VM **305** before storage. Source storage device **105-1** might decrypt, but not decompress, the data for VM **305** before sending the data for VM **305** as part of migrating VM **305** from source machine **105-1** to destination machine **105-2**.

[0132] In addition, while the above discussion addresses compression, encoding, encryption, and error correcting codes as possible manipulations of the data for VM **305**, source storage device **125-1** may make other manipulations of the data for VM **305** as well. These other manipulations may be undone before sending the data for VM **305** or not, as desired and appropriate for storage devices **125-1** and **125-2**.

[0133] As mentioned above, storage devices **125-1** and **125-2** may each act as source and/or destination, depending on what VM is being migrated and between what machines. Thus, storage devices **125-1** and **125-2** are each capable of performing the operations attributed above to source storage device **125-1** and destination storage device **125-2** separately. That is, storage devices **125-1** and **125-2** may each be notified by hypervisors **340-1** or **340-2** that VM migration is to occur, either may collect or set storage state **410** of FIG. 4 (and generate log pages or reports as appropriate), either may send or receive data for VM **305**, and either may provide information about deltas to either storage state **410** of FIG. 4 or the data for VM **305**. Thus, the operations attributed to either source storage device **125-1** or destination storage device **125-2** may both be performed by both source storage device **125-1** or destination storage device **125-2**.

[0134] The above discussion focuses only on operations relating to migration of information relating to VM **305** as stored on source storage device **125-1**. VM migration may also involve transfer of information relating to applications **310** of FIG. 3, the state of system software stack **315** of FIG. 3, the state of operating system **320** of FIG. 3, the state of memory **325** of FIG. 3, the state of network driver **330** of FIG. 3, and the state of storage driver **335** of FIG. 3. This information may not be available to source storage device **125-1**, and therefore may be transferred by hypervisor **340-1** of source machine **105-1** to hypervisor **340-2** of destination machine **105-2** in addition to the above-described process.

[0135] While the above process is described with reference to VM migration, embodiments of the inventive concept support using the VM migration operations of storage device **125** for other purposes. For example, as mentioned above, hypervisor **340-1** of source machine **105-1** may use a snapshot to transfer the data for VM **305** to destination storage device **125-2**. Hypervisor **340-1** of source machine **105-1** (or indeed, any process on source machine **105-1**, even if source machine **105-1** is not supporting VMs) may request a snapshot be taken of data on source storage device **125-1**. Hypervisor **340-1** of source machine **105-1**, or any process on source machine **105-1**, may

then request source storage device **125-1** in the state as though a VM is to be migrated. Source storage device **125-1** may then track changes to data stored on storage device **125-1**. Hypervisor **340-1** of source machine **105-1**, or any process on source machine **105-1**, may then request the changes relative to the snapshot, much the same as hypervisor **340-1** of source machine **105-1** may request the deltas to the data for VM **305** on source storage device **125-1**. To that end, VM Migration State Monitor and Capture Module **415** of FIG. **4** may also be called a delta module (although the functionality of a delta module may be less than the functionality of the VM Migration State Monitor and Capture Module **415** of FIG. **4**, since source storage device **125-1** may not capture or track changes in storage state **410** of FIG. **4**, or may not be used to configure source storage device **125-1** in a particular state).

[0136] The above discussion includes references to various commands that may be sent to storage devices **125-1** and **125-2**. The following provides some additional detail regarding these commands.

[0137] NVMe commands include commands called Set Feature and Get Feature. These commands enable machines **105-1** and **105-2** to set (using Set Feature) or read (using Get Feature) the values of certain features of NVMe controller **350** of FIG. **3** within storage devices **125-1** and **125-2**.

Examples of such features may include command submission queue arbitration parameters, power management settings, temperature thresholds, numbers of queues, and interrupt coalescing configuration. To these features, new features, which may be labeled VM_MIGRATION, VM_STORAGE_STATE, VM_DIRTY_STATE, and VM_DIRTY_LBA_RANGES may be added.

[0138] Using the VM_MIGRATION feature, machines **105-1** and **105-2** may set or read certain parameters, such as those labeled VM_MIGRATION_ACTIVE, LBA_WRITTEN_THRESHOLD, and LBA_WRITTEN_THRESHOLD_EXCEEDED above.

[0139] VM_MIGRATION_ACTIVE may be used to inform NVMe controller **350** of FIG. **3** (and therefore storage devices **125-1** and **125-2**) that VM **305** is being migrated from source machine **105-1** to destination machine **105-2**. VM_MIGRATION_ACTIVE may also be read to determine whether NVMe controller **350** of FIG. **3** is currently actively assisting in migrating VM **305**.

[0140] LBA_WRITTEN_THRESHOLD may be used to inform NVMe controller **350** of FIG. **3** of how many LBA ranges to record in a log (as data is written to source storage device **125-1**). If more LBAs are written to source storage device **125-1** than the value of LBA_WRITTEN_THRESHOLD, storage device **125-1** may set the LBA_WRITTEN_THRESHOLD_EXCEEDED flag. Hypervisor **340-1** of source machine **105-1** may read the LBA_WRITTEN_THRESHOLD_EXCEEDED flag. If this flag is set, hypervisor **340-1** of source machine **105-1** may move the entire capacity allocated to VM **305** on source storage device **125-1** to destination storage device **125-2**. Otherwise, hypervisor **340-1** of source machine **105-1** may access a log page or other report that includes the LBAs written by VM **305** of FIG. **1**, and then readjust those LBAs from source storage device **125-1** to write that data to destination storage device **125-2**.

[0141] Two things are worth noting about LBA_WRITTEN_THRESHOLD and LBA_WRITTEN_THRESHOLD_EXCEEDED. First, while the parameters refer to “LBAs”, the parameters may be used to track data written using other methods. For example, in key-value storage devices, these parameters may be used to track how many keys have been written to source storage device **125-1** for VM **305** rather than LBAs. (The parameters may be renamed to KEY_WRITTEN_THRESHOLD and KEY_WRITTEN_THRESHOLD_EXCEEDED or some other comparable labels if desired: for purposes of this discussion, reference to LBAs should be understood as including references to other identifiers for data written to source storage device **125-1**.)

[0142] Second, since tracking how much data has been written to source storage device **125-1** may be meaningful over the life of the VM and not just since VM migration has started, hypervisor **340-1** of source machine **105-1** may set a value for the LBA_WRITTEN_THRESHOLD parameter

when VM **305** is first instantiated on source machine **105-1**. Thus, source storage device **125-1** may track the number of logical identifiers written to source storage device **125-1** over the life of VM **305** (or at least as long as VM **305** is instantiated on source machine **105-1**), rather than just since VM migration began.

[0143] Using the VM_STORAGE_STATE feature, machines **105-1** and **105-2** may set or read information about storage state **410** of FIG. 4. Storage state **410** of FIG. 4 is discussed further with reference to FIG. 6 below.

[0144] Using the VM_DIRTY_STATE feature, source machine **105-1** may read what deltas (changes) have been made to storage state **410** of FIG. 4 in source storage device **125-1** during the brownout phase. These deltas may be transferred to destination storage device **125-1** during the blackout phase. For VM_DIRTY_STATE to track this information, source machine **105-1** may set a flag indicating that any further changes to storage state **410** of FIG. 4 should be tracked by source storage device **125-1**. Alternatively, source storage device **125-1** may assume that any changes made to storage state **410** of FIG. 4 after source machine **105-1** has most recently read storage state **410** of FIG. 4 should be tracked by VM_DIRTY_STATE. Note that VM_DIRTY_STATE does not necessarily need a counterpart for setting the state of destination storage device **125-2**: changes to storage state **410** of FIG. 4 in source storage device **125-1** may be written to destination storage device **125-2** using the VM_STORAGE_STATE feature. But embodiments of the inventive concept may include a special command to set the VM_DIRTY_STATE feature to write such changes to storage state **410** of FIG. 4 for destination storage device **125**.

[0145] Using the VM_DIRTY_LBA_RANGES feature, machine **105-1** may read the list of LBAs (or other logical identifiers) that have been changed (either written or deleted) since hypervisor **305-1** of source machine **105-1** transferred the data for VM **305** from source storage device **125-1** to destination storage device **125-2** during the brownout phase. These deltas may be transferred to destination storage device **125-1** during the blackout phase. For VM_DIRTY_LBA_RANGES to track this information, source machine **105-1** may set a flag indicating that any further changes to the data for VM **305** should be tracked by source storage device **125-1**. Alternatively, source storage device **125-1** may assume that any changes made to the data for VM **305** after source machine **105-1** has transferred the data for VM **305** to destination storage device **125-2** should be tracked by VM_DIRTY_LBA_RANGES. Note that VM_DIRTY_LBA_RANGES does not necessarily need a counterpart for storing the data for VM **305** in destination storage device **125-2**, since the data may be written to destination storage device **125-2** using conventional write requests. But embodiments of the inventive concept may include a special command to write changes to the data for VM **305** to destination storage device **125**.

[0146] The NVMe specification includes an Identify data structure. This Identify data structure may be used indicate VM Live Migration Assist support to the hypervisors **340-1** and **340-2**. The unused fields of the Identify structure may be used for this purpose.

[0147] While the above describes how existing NVMe commands may be augmented to support VM migration, embodiments of the inventive concept may also use other commands, including custom vendor-defined commands, to implement support for VM migration.

[0148] FIG. 6 shows details of state information that may be maintained by storage device **125** of FIG. 1. In FIG. 6, storage state **410** is shown. Storage state **410** may include various information, such as transport configuration information **605**, interface configuration information **610**, I/O activity **615**, and translation table **620**. Transport configuration information **605** may include information about how the transport protocol is configured. For example, if storage device **125** of FIG. 1 is communicates with machine **105** of FIG. 1 over a Peripheral Component Interconnect Express (PCIe) connection, then transport configuration information **605** may include information about the PCIe configuration of storage device **125** of FIG. 1. Examples of information that may be considered transport configuration information **605** may include various PCIe capability settings, such as message signaled interrupts (MSI and MSI-X), advanced error reporting (AER), power

management, etc., and configuration space settings, such as interrupt information, maximum payload size, etc.

[0149] Interface configuration information **610** may include information about the configuration of the interface in storage device **125** of FIG. **1**. For example, if storage device **125** of FIG. **1** uses NVMe to communicate with machine **105** of FIG. **1**, then interface configuration information **610** may include information about the NVMe configuration of storage device **125** of FIG. **1**. Examples of information that may be considered interface configuration information **610** may include memory page size, doorbell stride, timeout values, submission queue arbitration settings, power management, number of queues, asynchronous event notification, interrupt coalescing, atomicity, namespaces created, security settings, log pages created so far, etc.

[0150] I/O activity **615** may include information about I/O activity for storage device **125** of FIG. **1**. This information may include VM data writes in terms of LBA ranges, and whether the number of VM data writes exceeds a threshold (such as the VM_DIRTY_LBA_RANGES and LBA_WRITTEN_THRESHOLD as discussed above).

[0151] Translation table **620** may include information about the translation table(s) used to translate logical identifiers to physical addresses in storage device **125** of FIG. **1**. Translation table **620** may also include other information, such as flash media block usage, write/read cache state, and settings for the NVMe controller **350** of FIG. **3** (and in turn VM **305**).

[0152] Some or all of the information in storage state **410** may be used by destination storage device **125-2** of FIG. **5** to reproduce the state of source storage device **125-1** of FIG. **5**. For example, transport configuration information **605** may be used to configure the transport protocol of destination storage device **125-2** of FIG. **5** to match that of source storage device **125-1** of FIG. **5**, interface configuration information **610** may be used to configure the interface of destination storage device **125-2** of FIG. **5** to match that of source storage device **125-1** of FIG. **5**, I/O activity **615** may be used to information destination storage device **125-2** of FIG. **5** about the I/O activity of VM **305** of FIG. **3** as known by source storage device **125-1** of FIG. **5**, and translation table **620** may be used to arrange the translation table and cache of storage device **125-2** of FIG. **5** to match, as much as possible, the configuration of source storage device **125-1** of FIG. **5**.

[0153] FIG. **7** shows VM Migration State Monitor and Capture Module **415** of FIG. **4** determining whether a threshold number of logical block addresses (LBAs) have been written. As discussed above, storage device **125** of FIG. **1** may track number **705** of LBAs written by VM **305** of FIG. **3**. VM Migration State Monitor and Capture Module **415** may compare number **705** with threshold **710**: if number **705** exceeds threshold **710**, then VM Migration State Monitor and Capture Module **415** may set a flag (such as LBA_WRITTEN_THRESHOLD_EXCEEDED) indicating that enough data has been written for VM **305** of FIG. **3** to storage device **125** of FIG. **1** that the entire capacity of source storage device **125-1** of FIG. **5** allocated to VM **305** to destination storage device **125-2** of FIG. **5**. Otherwise, storage device **125** of FIG. **1** may track the individual LBA addresses (or ranges thereof) that have been written to storage device **125** of FIG. **1** for VM **305** of FIG. **3**.

[0154] While FIG. **7** shows storage device **125** of FIG. **1** tracking number **705** of LBAs written by VM **305** of FIG. **3**, embodiments of the inventive concept may track other criteria. For example, embodiments of the inventive concept may track the space occupied on storage device **125** of FIG. **1**, a number of logical identifiers (which may differ from LBAs: for example, keys used in a key-value storage device), or a range of LBAs written by VM **305** of FIG. **3**. If other information may be tracked by storage device **125** of FIG. **1**, then VM Migration State Monitor and Capture Module **415** may use other thresholds for comparison.

[0155] FIG. **8** shows VM Migration State Monitor and Capture Module **415** of FIG. **4** generating a log page for the state of source storage device **125-1** of FIG. **5**. In FIG. **8**, VM Migration State Monitor and Capture Module **415** may access storage state **410** (or a portion thereof), and generate log page or report **805** therefrom. Log page or report **805** may take any desired form and/or format, and should be suitable for use by hypervisor **340-2** of FIG. **5** of destination storage device **125-2** of

FIG. 5 to configure destination storage device **125-2** accordingly.

[0156] FIGS. **9A-9B** show a flowchart of an example procedure for storage device **125** of FIG. **1** to support live migration of a VM from machine **105** of FIG. **1**, according to embodiments of the inventive concept. In FIG. **9A**, at block **905**, source storage device **125-1** of FIG. **5** may receive a request from hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5** that VM **305** of FIG. **5** is to be migrated (from source machine **105-1** of FIG. **5** to destination machine **105-2** of FIG. **5**, although source storage device **125-1** of FIG. **5** may not be aware of destination machine **105-2** of FIG. **5**). At block **910**, source storage device **125-1** of FIG. **5** may modify cache policy **425** and/or prefetch policy **425** of FIG. **4** to favor keeping and/or moving data for VM **305** of FIG. **5** into cache **420** of FIG. **4**.

[0157] At block **915**, source storage device **125-1** of FIG. **5** may receive a request from hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5** for storage state **410** of FIG. **4** for VM **305** of FIG. **5**. In response, at block **920**, source storage device **125-1** of FIG. **5** may send storage state **410** of FIG. **4**, to any of hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5**, hypervisor **340-2** of FIG. **5** of destination machine **105-2** of FIG. **5**, or destination storage device **125-2** of FIG. **5**.

[0158] At block **925**, source storage device **125-1** of FIG. **5** may receive a request from hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5**, requesting the data for VM **305** of FIG. **5**. At block **910**, source storage device **125-1** of FIG. **5** may decompress, decode, decrypt, and/or remove error correcting codes from the data for VM **305** of FIG. **5**. Note that block **930** may be omitted if the data for VM **305** of FIG. **5** is to be sent compressed, encoded, encrypted, and with error correcting codes, as shown by dashed line **935**. In addition, at block **930** source storage device **125-1** of FIG. **1** may perform some operations but not others, such as decrypting but not decompressing the data for VM **305** of FIG. **5**, depending on whether destination storage device **125-2** of FIG. **5** may be able to process manipulated data for VM **305** of FIG. **1**.

[0159] At block **940** (FIG. **9B**), source storage device **125-1** of FIG. **5** may send the data for VM **305** of FIG. **5**. As discussed above, source storage device **125-1** of FIG. **5** may send the data for VM **305** of FIG. **5** to any of hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5**, hypervisor **340-2** of FIG. **5** of destination storage device **105-2** of FIG. **5**, or destination storage device **125-2** of FIG. **5**.

[0160] At block **945**, source storage device **125-1** of FIG. **5** may track any deltas (that is, changes) to storage state **410** of FIG. **5**. At block **950**, source storage device **125-1** of FIG. **5** may track any deltas (that is, changes) to the data for VM **305** of FIG. **5**. Note that blocks **945** and **950** may be performed in parallel rather than being sequential operations.

[0161] At block **955**, source storage device **125-1** of FIG. **5** may receive from hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5** a request for the deltas to storage state **410** of FIG. **4**. At block **960**, source storage device **125-1** of FIG. **5** may send the deltas to storage state **410** of FIG. **4**. As discussed above, source storage device **125-1** of FIG. **5** may send the deltas to storage state **410** of FIG. **4** to any of hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5**, hypervisor **340-2** of FIG. **5** of destination storage device **105-2** of FIG. **5**, or destination storage device **125-2** of FIG. **5**.

[0162] At block **965**, source storage device **125-1** of FIG. **5** may receive from hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5** a request for the deltas to the data for VM **305** of FIG. **5**. At block **970**, source storage device **125-1** of FIG. **5** may send the deltas to the data for VM **305** of FIG. **5**. As discussed above, source storage device **125-1** of FIG. **5** may send the deltas to the data for VM **305** of FIG. **5** to any of hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5**, hypervisor **340-2** of FIG. **5** of destination storage device **105-2** of FIG. **5**, or destination storage device **125-2** of FIG. **5**. In addition, while not explicitly shown in FIG. **9B**, source storage device **125-1** of FIG. **5** may decompress, decode, decrypt, and/or remove error correcting codes, as appropriate, the deltas to the data for VM **305** of FIG. **5** before sending the deltas to the data for VM **305** of FIG. **5**.

[0163] Finally, at block **975**, source storage device **125-1** of FIG. **5** may receive a message indicating that the migration of VM **305** of FIG. **5** from source machine **105-1** of FIG. **5** to destination machine **105-2** of FIG. **5** is complete, at which point source storage device **125-1** of FIG. **5** may stop assisting with the VM migration. Block **975** may be omitted, as shown by dashed line **980**, since once VM **305** of FIG. **5** has been migrated from source machine **105-1** of FIG. **5** to destination machine **105-2** of FIG. **5**, any data for VM **305** of FIG. **5** may be removed from source storage device **105-1** of FIG. **5**, which would effectively end storage device **125-1** of FIG. **5** assisting in the migration of VM **305** of FIG. **5** away from source storage device **125-1** of FIG. **1**.

[0164] FIG. **10** shows a flowchart of an example procedure for VM Migration State Monitor and Capture Module **415** of FIG. **4** to determine whether a threshold number of LBAs have been written, according to an embodiment of the inventive concept. In FIG. **10**, at block **1005** source storage device **125-1** of FIG. **1** may track number **705** of LBAs written for VM **305** of FIG. **5**. As discussed above, this tracking may be by total occupied storage space, number of LBAs (or other logical identifiers) written, the range of LBAs written for VM **305** of FIG. **5**, or other criteria. At block **1010**, VM Migration State Monitor and Capture Module **415** of FIG. **4** may compare number **705** of LBAs written for VM **305** of FIG. **5** with threshold **710**. If number **705** of LBAs written for VM **305** of FIG. **5** exceeds threshold **710**, then at block **1015** VM Migration State Monitor and Capture Module **415** of FIG. **4** may set a flag indicating that enough data has been written for VM **305** of FIG. **5** that the entire capacity of source storage device **125-1** of FIG. **5** allocated to VM **305** of FIG. **5** may be transferred to destination storage device **125-2** of FIG. **5** rather than a select set of data for VM **305** of FIG. **5**.

[0165] FIG. **11** shows a flowchart of an example procedure for VM Migration State Monitor and Capture Module **415** of FIG. **4** to generate and deliver a report of storage state **410** of FIG. **4** for the VM, according to an embodiment of the inventive concept. In FIG. **11**, at block **1105**, VM Migration State Monitor and Capture Module **415** of FIG. **4** may (at the request of hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **1**) generate a log page or other report regarding storage state **410** of FIG. **4** for VM **305** of FIG. **5**. This log page or other report may include, for example, transport configuration information **605** of FIG. **6**, interface configuration information **610** of FIG. **6**, I/O activity **615** of FIG. **6**, and translation table **620** of FIG. **6**, or the log page or report may include just a subset of these data. At block **1110**, VM Migration State Monitor and Capture Module **415** of FIG. **4** may store the log page or other report in storage somewhere in source storage device **125-1** of FIG. **5**, so that (at block **1115**) hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5** may read the log page or other report from source storage device **125-1** of FIG. **5**. Alternatively, at block **1120**, VM Migration State Monitor and Capture Module **415** of FIG. **4** may send the log page or other report to hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **6** (or to some other destination, such as hypervisor **340-2** of FIG. **5** of destination machine **105-2** of FIG. **5** or destination storage device **125-2** of FIG. **5**) in a message.

[0166] FIGS. **12A-12B** show a flowchart of an example procedure for storage device **125** of FIG. **1** to support live migration of a VM to machine **105** of FIG. **1**, according to embodiments of the inventive concept. In FIG. **12A**, at block **1205**, destination storage device **125-2** of FIG. **5** may receive a request from hypervisor **340-2** of FIG. **5** of destination machine **105-2** of FIG. **5** to assist in the migration of VM **305** of FIG. **5** from source machine **105-1** of FIG. **5** to destination machine **105-2** of FIG. **5**. At block **1210**, destination storage device **125-2** of FIG. **5** may receive storage state **410** of FIG. **4**. Destination storage device **125-2** of FIG. **5** may receive storage state **410** of FIG. **4** from any of hypervisor **340-1** of FIG. **5** of source machine **105-1** of FIG. **5**, hypervisor **340-2** of FIG. **5** of destination machine **105-2** of FIG. **5**, or source storage device **125-1** of FIG. **5**. At block **1215**, destination storage device **125-2** of FIG. **5** may configure itself appropriately for VM **305** of FIG. **5** based on the information received in storage state **410** of FIG. **4**.

[0167] At block **1220**, destination storage device **125-2** of FIG. **5** may receive data for VM **305** of FIG. **5**. Destination storage device **125-2** of FIG. **5** may receive data for VM **305** of FIG. **5** from

any of hypervisor **340-1** of FIG. 5 of source machine **105-1** of FIG. 5, hypervisor **340-2** of FIG. 5 of destination machine **105-2** of FIG. 5, or source storage device **125-1** of FIG. 5. At block **1225**, destination storage device **125-2** of FIG. 5 may compress, encode, encrypt, and/or introduce error correcting codes to the data for VM **305** of FIG. 5 as appropriate. Block **1225** may be omitted, as shown by dashed line **1230**. Then, at block **1235**, destination storage device **125-2** of FIG. 5 may store the data for VM **305** of FIG. 5.

[0168] At block **1240** (FIG. 12B), destination storage device **125-2** of FIG. 5 may receive deltas (changes) in storage state **410** of FIG. 4. Destination storage device **125-2** of FIG. 5 may receive the deltas for storage state **410** of FIG. 4 from any of hypervisor **340-1** of FIG. 5 of source machine **105-1** of FIG. 5, hypervisor **340-2** of FIG. 5 of destination machine **105-2** of FIG. 5, or source storage device **125-1** of FIG. 5. At block **1245**, destination storage device **125-2** of FIG. 2 may update its storage state consistent with the deltas to storage state **410** of FIG. 4.

[0169] At block **1250**, destination storage device **125-2** of FIG. 5 may receive deltas (changes) to the data for VM **305** of FIG. 5. Destination storage device **125-2** of FIG. 5 may receive the deltas to the data for VM **305** of FIG. 5 from any of hypervisor **340-1** of FIG. 5 of source machine **105-1** of FIG. 5, hypervisor **340-2** of FIG. 5 of destination machine **105-2** of FIG. 5, or source storage device **125-1** of FIG. 5. At block **1255**, destination storage device **125-2** of FIG. 5 may update the data for VM **305** of FIG. 5. In addition, while not explicitly shown in FIG. 12B, destination storage device **125-2** of FIG. 5 may compress, encode, encrypt, and/or introduce error correcting codes to, as appropriate, the deltas to the data for VM **305** of FIG. 5 before storing the deltas to the data for VM **305** of FIG. 5.

[0170] Finally, at block **1260**, destination storage device **125-2** of FIG. 5 may receive a message from hypervisor **340-2** of FIG. 5 of destination machine **105-2** of FIG. 5 that the migration of VM **305** of FIG. 5 from source machine **105-1** of FIG. 5 to destination machine **105-2** of FIG. 5 is complete, at which point destination storage device **125-2** of FIG. 5 may stop assisting with the VM migration. Unlike source storage device **125-1** of FIG. 5 (which may not receive a message that migration of VM **305** of FIG. 5 from source storage device **125-1** of FIG. 5 to destination storage device **125-2** of FIG. 5 is complete), destination storage device **125-2** of FIG. 5 should receive such a message, since destination storage device **125-2** of FIG. 5 may support the operation of VM **305** of FIG. 5 going forward.

[0171] In FIGS. 9A-12B, some embodiments of the inventive concept are shown. But a person skilled in the art will recognize that other embodiments of the inventive concept are also possible, by changing the order of the blocks, by omitting blocks, or by including links not shown in the drawings, irrespective of any elements that may specifically be omitted. All such variations of the flowcharts are considered to be embodiments of the inventive concept, whether expressly described or not.

[0172] Embodiments of the inventive concept include technical advantages over conventional implementations. By using a storage device to assist in the migration of a VM, the load on the hypervisor managing the migration of the VM may be reduced. By reducing the load on the hypervisor, the hypervisor may be freed to handle other tasks, such as managing the execution of other VMs, selecting other VMs for migration to other machines, and the like. Embodiments of the inventive concept also result in a storage device that is aware of VM migration, and which may shift resources to benefit VM migration. For example, a cache in the storage device may be used to store data for the VM, since accessing data from the cache may be faster than accessing data from the underlying storage medium of the storage device.

[0173] The following discussion is intended to provide a brief, general description of a suitable machine or machines in which certain aspects of the inventive concept can be implemented. The machine or machines can be controlled, at least in part, by input from conventional input devices, such as keyboards, mice, etc., as well as by directives received from another machine, interaction with a virtual reality (VR) environment, biometric feedback, or other input signal. As used herein,

the term “machine” is intended to broadly encompass a single machine, a virtual machine, or a system of communicatively coupled machines, virtual machines, or devices operating together. Exemplary machines include computing devices such as personal computers, workstations, servers, portable computers, handheld devices, telephones, tablets, etc., as well as transportation devices, such as private or public transportation, e.g., automobiles, trains, cabs, etc.

[0174] The machine or machines can include embedded controllers, such as programmable or non-programmable logic devices or arrays, Application Specific Integrated Circuits (ASICs), embedded computers, smart cards, and the like. The machine or machines can utilize one or more connections to one or more remote machines, such as through a network interface, modem, or other communicative coupling. Machines can be interconnected by way of a physical and/or logical network, such as an intranet, the Internet, local area networks, wide area networks, etc. One skilled in the art will appreciate that network communication can utilize various wired and/or wireless short range or long range carriers and protocols, including radio frequency (RF), satellite, microwave, Institute of Electrical and Electronics Engineers (IEEE) 802.11, Bluetooth®, optical, infrared, cable, laser, etc.

[0175] Embodiments of the present inventive concept can be described by reference to or in conjunction with associated data including functions, procedures, data structures, application programs, etc. which when accessed by a machine results in the machine performing tasks or defining abstract data types or low-level hardware contexts. Associated data can be stored in, for example, the volatile and/or non-volatile memory, e.g., RAM, ROM, etc., or in other storage devices and their associated storage media, including hard-drives, floppy-disks, optical storage, tapes, flash memory, memory sticks, digital video disks, biological storage, etc. Associated data can be delivered over transmission environments, including the physical and/or logical network, in the form of packets, serial data, parallel data, propagated signals, etc., and can be used in a compressed or encrypted format. Associated data can be used in a distributed environment, and stored locally and/or remotely for machine access.

[0176] Embodiments of the inventive concept can include a tangible, non-transitory machine-readable medium comprising instructions executable by one or more processors, the instructions comprising instructions to perform the elements of the inventive concepts as described herein.

[0177] Having described and illustrated the principles of the inventive concept with reference to illustrated embodiments, it will be recognized that the illustrated embodiments can be modified in arrangement and detail without departing from such principles, and can be combined in any desired manner. And, although the foregoing discussion has focused on particular embodiments, other configurations are contemplated. In particular, even though expressions such as “according to an embodiment of the inventive concept” or the like are used herein, these phrases are meant to generally reference embodiment possibilities, and are not intended to limit the inventive concept to particular embodiment configurations. As used herein, these terms can reference the same or different embodiments that are combinable into other embodiments.

[0178] The foregoing illustrative embodiments are not to be construed as limiting the inventive concept thereof. Although a few embodiments have been described, those skilled in the art will readily appreciate that many modifications are possible to those embodiments without materially departing from the novel teachings and advantages of the present disclosure. Accordingly, all such modifications are intended to be included within the scope of this inventive concept as defined in the claims.

[0179] Embodiments of the inventive concept can extend to the following statements, without limitation:

[0180] Statement 1. An embodiment of the inventive concept includes a storage device, comprising: [0181] at least one controller for a virtual machine (VM), the VM on a source host; [0182] first storage for a storage data for the VM; [0183] second storage for a storage state for the VM; [0184] a storage device controller to process at least one read request received from the

controller for the VM using the first storage and to process at least one write request received from the controller for the VM using the first storage; and [0185] a VM migration state monitor and capture module to assist in the migration of the VM from the source host to a destination host.

[0186] Statement 2. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the storage device includes a Solid State Drive (SSD).

[0187] Statement 3. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the storage state includes at least one of a transport configuration information for the controller for the VM, an interface configuration information for the controller for the VM, an input/output (I/O) activity for the VM, or a table to translate a logical identifier used by the VM to a physical address used by the storage device.

[0188] Statement 4. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the VM migration state monitor and capture module tracks first changes in a storage state for the VM in the storage device and tracks second changes in the storage data for the VM in the storage device.

[0189] Statement 5. An embodiment of the inventive concept includes the storage device according to statement 4, wherein the VM migration state monitor and capture module sends the first changes and the second changes to at least one of a hypervisor on the source host and a second storage device on the destination host.

[0190] Statement 6. An embodiment of the inventive concept includes the storage device according to statement 4, wherein the VM migration state monitor and capture module sends the storage state and the storage data for the VM to at least one of a hypervisor on the source host and a second storage device on the destination host.

[0191] Statement 7. An embodiment of the inventive concept includes the storage device according to statement 6, wherein the storage data is subject to at least one of compression, encoding, or encryption.

[0192] Statement 8. An embodiment of the inventive concept includes the storage device according to statement 7, wherein the VM migration state monitor and capture module sends the storage state and the storage data for the VM to at least one of a hypervisor on the source host and a second storage device on the destination host after at least one of decompressing, decoding, or decrypting the storage data.

[0193] Statement 9. An embodiment of the inventive concept includes the storage device according to statement 4, wherein the VM migration state monitor and capture module sets a flag based at least in part on an amount of data written by the VM to the storage device exceeding a threshold.

[0194] Statement 10. An embodiment of the inventive concept includes the storage device according to statement 9, wherein a hypervisor on the source uses the flag to select between transferring all storage data for the VM from the storage device or a subset of the storage data for the VM from the storage device.

[0195] Statement 11. An embodiment of the inventive concept includes the storage device according to statement 4, wherein the VM migration state monitor and capture module generates a report including the storage state.

[0196] Statement 12. An embodiment of the inventive concept includes the storage device according to statement 1, further comprising a cache and a cache policy.

[0197] Statement 13. An embodiment of the inventive concept includes the storage device according to statement 12, wherein the cache policy retains a cache line in the cache until the cache line is sent to at least one of a hypervisor on the source host or a second storage device on the destination host.

[0198] Statement 14. An embodiment of the inventive concept includes the storage device according to statement 1, further comprising a cache and a prefetch policy.

[0199] Statement 15. An embodiment of the inventive concept includes the storage device according to statement 14, wherein the prefetch policy transfers storage data for the VM from the

first storage to the cache to expedite sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0200] Statement 16. An embodiment of the inventive concept includes the storage device according to statement 15, wherein the storage device performs at least one of decompressing, decoding, or decrypting the storage data transferred from the first storage to the cache according to the prefetch policy.

[0201] Statement 17. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the at least one controller includes a Non-Volatile Memory Express (NVMe) controller.

[0202] Statement 18. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the at least one controller is associated with one of a physical function (PF) or a virtual function (VF) exposed for the storage device.

[0203] Statement 19. An embodiment of the inventive concept includes the storage device according to statement 18, wherein the at least one controller is associated with a unique one of the PF or the VF.

[0204] Statement 20. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the at least one controller includes at least one of a Field Programmable Gate Array (FPGA), an Application-Specific Integrated Circuit (ASIC), a central processing unit (CPU), a graphics processing unit (GPU), a general purpose GPU (GPGPU), or software.

[0205] Statement 21. An embodiment of the inventive concept includes the storage device according to statement 1, wherein the storage device controller includes the at least one controller.

[0206] Statement 22. An embodiment of the inventive concept includes a storage device, comprising: [0207] first storage for a storage data; [0208] a storage device controller to process at least one read request received from a host using the first storage and to process at least one write request received from the host using the first storage; and [0209] a delta module to identify changes in the storage data in the first storage relative to a snapshot.

[0210] Statement 23. An embodiment of the inventive concept includes the storage device according to statement 22, wherein the delta module sends the changes in the storage data to the host.

[0211] Statement 24. An embodiment of the inventive concept includes a method, comprising: [0212] receiving, at a storage device on a source host, a command to assist in the migration of a virtual machine (VM) from the source host to a destination host; [0213] tracking first changes in a storage state for the VM in the storage device; [0214] tracking second changes in a storage data for the VM in the storage device; [0215] sending the first changes in the storage state for the VM from the storage device; and [0216] sending the second changes in the storage data for the VM from the storage device.

[0217] Statement 25. An embodiment of the inventive concept includes the method according to statement 24, wherein the storage device includes a Solid State Drive (SSD).

[0218] Statement 26. An embodiment of the inventive concept includes the method according to statement 24, wherein: [0219] sending the first changes in the storage state from the storage device includes sending the first changes in the storage state from the storage device to at least one of a hypervisor on the source host or a second storage device on the destination host; and [0220] sending the second changes in the storage data for the from the storage device includes sending the second changes in the storage data for the from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0221] Statement 27. An embodiment of the inventive concept includes the method according to statement 24, further comprising: [0222] sending the storage state for the VM from the storage device; and [0223] sending the storage data for the VM from the storage device.

[0224] Statement 28. An embodiment of the inventive concept includes the method according to statement 27, wherein: [0225] sending the storage state for the VM from the storage device includes sending the storage state for the VM from the storage device to at least one of a hypervisor on the source host or a second storage device on the destination host; and [0226] sending the storage data for the VM from the storage device includes sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0227] Statement 29. An embodiment of the inventive concept includes the method according to statement 28, wherein the storage data is subject to at least one of compression, encoding, or encryption.

[0228] Statement 30. An embodiment of the inventive concept includes the method according to statement 29, wherein sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host includes: [0229] at least one of decompressing, decoding, or decrypting the storage data to produce non-coded storage data; and [0230] sending the non-coded storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0231] Statement 31. An embodiment of the inventive concept includes the method according to statement 27, wherein: [0232] the method further comprises receiving, at the storage device, a request for the storage state for the VM; and [0233] sending the storage state for the VM from the storage device includes sending the storage state for the VM from the storage device responsive to the request for the storage state for the VM.

[0234] Statement 32. An embodiment of the inventive concept includes the method according to statement 27, wherein: [0235] the method further comprises receiving, at the storage device, a request for the storage data for the VM; and [0236] sending the storage data for the VM from the storage device includes sending the storage data for the VM from the storage device responsive to the request for the storage data for the VM.

[0237] Statement 33. An embodiment of the inventive concept includes the method according to statement 27, further comprising: [0238] tracking an amount of data written by the VM to the storage device; and [0239] based at least in part on the amount of data written by the VM to the storage device exceeding a threshold, setting a flag in the storage device, [0240] wherein the source host may use the flag to select between transferring all storage data for the VM from the storage device or a subset of the storage data for the VM from the storage device.

[0241] Statement 34. An embodiment of the inventive concept includes the method according to statement 27, wherein sending the storage state for the VM from the storage device includes: [0242] generating a report including the storage state for the VM; and [0243] storing the report on the storage device, [0244] wherein a hypervisor on the source host may read the report from the storage device.

[0245] Statement 35. An embodiment of the inventive concept includes the method according to statement 24, wherein the storage state includes at least one of a transport configuration information for a controller for the VM, an interface configuration information for the controller for the VM, an input/output (I/O) activity for the VM, or a table to translate a logical identifier used by the VM to a physical address used by the storage device.

[0246] Statement 36. An embodiment of the inventive concept includes the method according to statement 24, further comprising adjusting at least one of a cache policy or a prefetch policy to assist in the migration of the VM from the source host to the destination host.

[0247] Statement 37. An embodiment of the inventive concept includes the method according to statement 36, wherein the cache policy retains a cache line in a cache until the cache line is sent to at least one of a hypervisor on the source host or a second storage device on the destination host; and [0248] the prefetch policy transfers storage data for the VM from storage to the cache to

expedite sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0249] Statement 38. An embodiment of the inventive concept includes the method according to statement 37, further comprising at least one of decompressing, decoding, or decrypting the storage data for the VM in the cache.

[0250] Statement 39. An embodiment of the inventive concept includes the method according to statement 24, further comprising receiving, at a storage device on a source host, a command that the migration of the virtual machine (VM) from the source host to a destination host is complete.

[0251] Statement 40. An embodiment of the inventive concept includes a method, comprising:

[0252] receiving, at a storage device on a destination host, a command to assist in the migration of a virtual machine (VM) from a source host to the destination host; [0253] receiving, at the storage device, changes in a storage state for the VM; [0254] updating the storage device consistent with the changes in the storage state for the VM; [0255] receiving, at the storage device, changes in a storage data for the VM; and [0256] updating the storage data for the VM on the storage device consistent with the changes in the storage data for the VM.

[0257] Statement 41. An embodiment of the inventive concept includes the method according to statement 40, further comprising receiving a second command from at the storage device that the VM is active.

[0258] Statement 42. An embodiment of the inventive concept includes the method according to statement 40, wherein the storage device includes a Solid State Drive (SSD).

[0259] Statement 43. An embodiment of the inventive concept includes the method according to statement 40, wherein: [0260] receiving, at the storage device, changes in a storage state for the VM includes receiving, at the storage device, changes in the storage state for the VM from at least one of a hypervisor on the destination host or a second storage device on the source host; [0261] receiving, at the storage device, changes in a storage data for the VM includes receiving, at the storage device, changes in the storage data for the VM from at least one of the hypervisor on the destination host or the second storage device on the source host.

[0262] Statement 44. An embodiment of the inventive concept includes the method according to statement 40, further comprising: [0263] receiving, at the storage device, the storage state for the VM; [0264] configuring the storage device consistent with the storage state for the VM; [0265] receiving, at the storage device, the storage data for the VM; and [0266] storing the storage data for the VM on the storage device.

[0267] Statement 45. An embodiment of the inventive concept includes the method according to statement 44, wherein: [0268] receiving, at the storage device, the storage state for the VM includes receiving, at the storage device, the storage state for the VM from at least one of a hypervisor on the destination host or a second storage device on the source host; [0269] receiving, at the storage device, the storage data for the VM includes receiving, at the storage device, the storage data for the VM from at least one of the hypervisor on the destination host or the second storage device on the source host.

[0270] Statement 46. An embodiment of the inventive concept includes the method according to statement 45, wherein the storage data is subject to at least one of compression, encoding, or encryption.

[0271] Statement 47. An embodiment of the inventive concept includes the method according to statement 46, wherein receiving, at the storage device, the storage data for the VM from at least one of the hypervisor on the destination host or the second storage device on the source host includes: [0272] receiving, at the storage device, non-coded data for the VM from at least one of the hypervisor on the source host or the second storage device on the destination host; and [0273] at least one of compressing, encoding, or encrypting the non-coded data to produce the storage data.

[0274] Statement 48. An embodiment of the inventive concept includes the method according to statement 40, wherein the storage state includes at least one of a transport configuration

information for a controller for the VM, an interface configuration information for the controller for the VM, an input/output (I/O) activity for the VM, or a table to translate a logical identifier used by the VM to a physical address used by the storage device.

[0275] Statement 49. An embodiment of the inventive concept includes an article, comprising a non-transitory storage medium, the non-transitory storage medium having stored thereon instructions that, when executed by a machine, result in: [0276] receiving, at a storage device on a source host, a command to assist in the migration of a virtual machine (VM) from the source host to a destination host; [0277] tracking first changes in a storage state for the VM in the storage device; [0278] tracking second changes in a storage data for the VM in the storage device; [0279] sending the first changes in the storage state for the VM from the storage device; and [0280] sending the second changes in the storage data for the VM from the storage device.

[0281] Statement 50. An embodiment of the inventive concept includes the article according to statement 49, wherein the storage device includes a Solid State Drive (SSD).

[0282] Statement 51. An embodiment of the inventive concept includes the article according to statement 49, wherein: [0283] sending the first changes in the storage state from the storage device includes sending the first changes in the storage state from the storage device to at least one of a hypervisor on the source host or a second storage device on the destination host; and [0284] sending the second changes in the storage data for the from the storage device includes sending the second changes in the storage data for the from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0285] Statement 52. An embodiment of the inventive concept includes the article according to statement 49, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in: [0286] sending the storage state for the VM from the storage device; and [0287] sending the storage data for the VM from the storage device.

[0288] Statement 53. An embodiment of the inventive concept includes the article according to statement 52, wherein: [0289] sending the storage state for the VM from the storage device includes sending the storage state for the VM from the storage device to at least one of a hypervisor on the source host or a second storage device on the destination host; and [0290] sending the storage data for the VM from the storage device includes sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0291] Statement 54. An embodiment of the inventive concept includes the article according to statement 53, wherein the storage data is subject to at least one of compression, encoding, or encryption.

[0292] Statement 55. An embodiment of the inventive concept includes the article according to statement 54, wherein sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host includes: [0293] at least one of decompressing, decoding, or decrypting the storage data to produce non-coded storage data; and [0294] sending the non-coded storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0295] Statement 56. An embodiment of the inventive concept includes the article according to statement 52, wherein: [0296] the non-transitory storage medium has stored thereon further instructions that, when executed by the machine, result in receiving, at the storage device, a request for the storage state for the VM; and [0297] sending the storage state for the VM from the storage device includes sending the storage state for the VM from the storage device responsive to the request for the storage state for the VM.

[0298] Statement 57. An embodiment of the inventive concept includes the article according to statement 52, wherein: [0299] the non-transitory storage medium has stored thereon further instructions that, when executed by the machine, result in receiving, at the storage device, a request

for the storage data for the VM; and [0300] sending the storage data for the VM from the storage device includes sending the storage data for the VM from the storage device responsive to the request for the storage data for the VM.

[0301] Statement 58. An embodiment of the inventive concept includes the article according to statement 52, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in: [0302] tracking an amount of data written by the VM to the storage device; and [0303] based at least in part on the amount of data written by the VM to the storage device exceeding a threshold, setting a flag in the storage device, [0304] wherein the source host may use the flag to select between transferring all storage data for the VM from the storage device or a subset of the storage data for the VM from the storage device.

[0305] Statement 59. An embodiment of the inventive concept includes the article according to statement 52, wherein sending the storage state for the VM from the storage device includes:

[0306] generating a report including the storage state for the VM; and [0307] storing the report on the storage device, [0308] wherein a hypervisor on the source host may read the report from the storage device.

[0309] Statement 60. An embodiment of the inventive concept includes the article according to statement 49, wherein the storage state includes at least one of a transport configuration information for a controller for the VM, an interface configuration information for the controller for the VM, an input/output (I/O) activity for the VM, or a table to translate a logical identifier used by the VM to a physical address used by the storage device.

[0310] Statement 61. An embodiment of the inventive concept includes the article according to statement 49, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in adjusting at least one of a cache policy or a prefetch policy to assist in the migration of the VM from the source host to the destination host.

[0311] Statement 62. An embodiment of the inventive concept includes the article according to statement 61, wherein the cache policy retains a cache line in a cache until the cache line is sent to at least one of a hypervisor on the source host or a second storage device on the destination host; and [0312] the prefetch policy transfers storage data for the VM from storage to the cache to expedite sending the storage data for the VM from the storage device to at least one of the hypervisor on the source host or the second storage device on the destination host.

[0313] Statement 63. An embodiment of the inventive concept includes the article according to statement 62, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in at least one of decompressing, decoding, or decrypting the storage data for the VM in the cache.

[0314] Statement 64. An embodiment of the inventive concept includes the article according to statement 49, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in receiving, at a storage device on a source host, a command that the migration of the virtual machine (VM) from the source host to a destination host is complete.

[0315] Statement 65. An embodiment of the inventive concept includes an article, comprising a non-transitory storage medium, the non-transitory storage medium having stored thereon instructions that, when executed by a machine, result in: [0316] receiving, at a storage device on a destination host, a command to assist in the migration of a virtual machine (VM) from a source host to the destination host; [0317] receiving, at the storage device, changes in a storage state for the VM; [0318] updating the storage device consistent with the changes in the storage state for the VM; [0319] receiving, at the storage device, changes in a storage data for the VM; and [0320] updating the storage data for the VM on the storage device consistent with the changes in the storage data for the VM.

[0321] Statement 66. An embodiment of the inventive concept includes the article according to statement 65, the non-transitory storage medium having stored thereon further instructions that,

when executed by the machine, result in receiving a second command from at the storage device that the VM is active.

[0322] Statement 67. An embodiment of the inventive concept includes the article according to statement 65, wherein the storage device includes a Solid State Drive (SSD).

[0323] Statement 68. An embodiment of the inventive concept includes the article according to statement 65, wherein: [0324] receiving, at the storage device, changes in a storage state for the VM includes receiving, at the storage device, changes in the storage state for the VM from at least one of a hypervisor on the destination host or a second storage device on the source host; [0325] receiving, at the storage device, changes in a storage data for the VM includes receiving, at the storage device, changes in the storage data for the VM from at least one of the hypervisor on the destination host or the second storage device on the source host.

[0326] Statement 69. An embodiment of the inventive concept includes the article according to statement 65, the non-transitory storage medium having stored thereon further instructions that, when executed by the machine, result in: [0327] receiving, at the storage device, the storage state for the VM; [0328] configuring the storage device consistent with the storage state for the VM; [0329] receiving, at the storage device, the storage data for the VM; and storing the storage data for the VM on the storage device.

[0330] Statement 70. An embodiment of the inventive concept includes the article according to statement 69, wherein: [0331] receiving, at the storage device, the storage state for the VM includes receiving, at the storage device, the storage state for the VM from at least one of a hypervisor on the destination host or a second storage device on the source host; [0332] receiving, at the storage device, the storage data for the VM includes receiving, at the storage device, the storage data for the VM from at least one of the hypervisor on the destination host or the second storage device on the source host.

[0333] Statement 71. An embodiment of the inventive concept includes the article according to statement 70, wherein the storage data is subject to at least one of compression, encoding, or encryption.

[0334] Statement 72. An embodiment of the inventive concept includes the article according to statement 71, wherein receiving, at the storage device, the storage data for the VM from at least one of the hypervisor on the destination host or the second storage device on the source host includes: [0335] receiving, at the storage device, non-coded data for the VM from at least one of the hypervisor on the source host or the second storage device on the destination host; and [0336] at least one of compressing, encoding, or encrypting the non-coded data to produce the storage data.

[0337] Statement 73. An embodiment of the inventive concept includes the article according to statement 65, wherein the storage state includes at least one of a transport configuration information for a controller for the VM, an interface configuration information for the controller for the VM, an input/output (I/O) activity for the VM, or a table to translate a logical identifier used by the VM to a physical address used by the storage device.

[0338] Consequently, in view of the wide variety of permutations to the embodiments described herein, this detailed description and accompanying material is intended to be illustrative only, and should not be taken as limiting the scope of the inventive concept. What is claimed as the inventive concept, therefore, is all such modifications as may come within the scope and spirit of the following claims and equivalents thereto.

Claims

1-10. (canceled)

11. A first storage device, comprising: a first storage medium; and a first storage device controller to access a data stored on the first storage medium; wherein the first storage device is configured to: receive a migration request from a host server, set an indicator to initiate migration in response

to receiving the migration request, obtain a state of the first storage device controller of a second storage device, and transfer the data and the state of the first storage device to a second storage device, the second storage device including a second storage medium and a second storage device controller, the second storage device associated with a destination server, and wherein the first storage device is associated with the host server.

12. The first storage device according to claim 11, wherein the data includes a virtual machine data for a virtual machine executing on the host server.

13. The first storage device according to claim 12, wherein the migration request is associated with a migration of the virtual machine from the host server to the destination server.

14. The first storage device according to claim 11, wherein the data includes a storage device configuration information associated with a configuration of the first storage device.

15. The first storage device according to claim 14, wherein the storage device configuration information includes the storage device configuration for a virtual machine executing on the host server.

16. The first storage device according to claim 11, wherein the first storage device is further configured to: identify changes in the data stored on the first storage device; and transfer the changes in the data to the second storage device based on the obtained state.

17. The first storage device according to claim 16, wherein the first storage device is further configured to receive a message that execution of a virtual machine on the host server has halted, the virtual machine associated with the data on the first storage device.

18. The first storage device according to claim 11, wherein the first storage device controller is configured to cache the data based at least in part on the migration request.

19. The first storage device according to claim 18, wherein the first storage device controller is configured to cache the data at a priority based at least in part on the migration request.

20. The first storage device according to claim 11, wherein the first storage device communicates directly with the second storage device to transfer the data to the second storage device.

21. A method, comprising: receiving, at a first storage device, a migration request from a host server, the first storage device including a first storage medium and a first storage device controller to access a data stored on the first storage medium, the first storage device associated with the host server; setting, at the first storage device, an indicator to initiate migration of the data in response to the first storage device receiving the migration request; obtaining, by the first storage device, a state of the first storage device controller; and transferring, by the first storage device, the data and the state of the first storage device to a second storage device, the second storage device including a second storage medium and a second storage device controller, the second storage device associated with a destination server.

22. The method according to claim 21, wherein the data includes a virtual machine data for a virtual machine executing on the host server.

23. The method according to claim 22, wherein receiving, at the first storage device, the migration request from the host server includes receiving, at the first storage device, the migration request from a hypervisor executing on the host server.

24. The method according to claim 23, wherein the migration request is associated with a migration of the virtual machine from the host server to the destination server.

25. The method according to claim 21, wherein the data includes a storage device configuration information associated with a configuration of the first storage device.

26. The method according to claim 25, wherein the storage device configuration information includes the storage device configuration for a virtual machine executing on the host server.

27. The method according to claim 21, further comprising: Identifying, by the first storage device, changes in the data stored on the first storage device; and transferring, from the first storage device, the changes in the data to the second storage device based on the obtained state.

28. The method according to claim 27, wherein transferring, from the first storage device, the

changes in the data to the second storage device based on the obtained state includes receiving, at the first storage device, a message that execution of a virtual machine on the host server has halted, the virtual machine associated with the data on the first storage device.

29. The method according to claim 21, further comprising caching, by the first storage device controller, the data based at least in part on the migration request.

30. The method according to claim 29, wherein caching, by the first storage device controller, the data based at least in part on the migration request includes caching, by the first storage device controller, the data at a priority based at least in part on the migration request.
