METHODS AND APPARATUS TO FACILITATE USE OF DYNAMIC COHERENCY MODELS FOR MEMORY OBJECTS

Abstract

Systems, apparatus, articles of manufacture, and methods are disclosed to facilitate use of dynamic coherency models for memory objects. An example apparatus includes interface circuitry; instructions; and at least one programmable circuit to be programmed by the instructions to: access a request identifying a memory range and a coherency model; and apply the coherency model to a memory address in the memory range based on the coherency model setting.

Inventors: **Guim Bernat; Francesc (Barcelona, ES), Gonzalez Pellicer; Edgar (Barcelona, ES), Mora Porta; Gaspar (Castellon de la Plana, ES), Moschiano; Violante (Avezzano, IT), Tagaya; Satoru (Barcelona, ES), Focht; Erich Ludwig (Stuttgart, DE), Vali; Tommaso (Sezze, IT)**

Applicant: **Openchip & Software Technologies S.L.** (Barcelona, ES)

Family ID: **95066063**

Appl. No.: **19/192032**

Filed: **April 28, 2025**

Foreign Application Priority Data

WO	PCT/EP2025/057026	Mar. 14, 2025

Related U.S. Application Data

parent WO continuation PCT/EP2025/057026 20250314 PENDING child US 19192032

Publication Classification

Int. Cl.: **G06F12/0831** (20160101)

U.S. Cl.:

CPC   **G06F12/0831** (20130101);

Background/Summary

RELATED APPLICATION(S) [0001] This patent arises from a continuation of International Patent Application No. PCT/EP2025/057026, which was filed on Mar. 14, 2025. Priority to International Patent Application No. PCT/EP2025/057026 is claimed. International Patent Application No. PCT/EP2025/057026 is incorporated herein by reference in its entirety.

FIELD OF THE DISCLOSURE
[0003] This disclosure relates generally to computing devices and, more particularly, to methods and apparatus to facilitate use of dynamic coherency models for memory objects.
BACKGROUND
[0004] Cache systems having multiple levels (L1 cache, L2 cache, etc.) are known. Cache systems typically include a persistent memory which retains large amounts of data but tends to require large amounts of time to access the data. One or more other layers of cache contain progressively smaller amounts of data with progressively faster access times. The cache closest to processor circuitry tends to be very fast access but stores a relatively small amount of data. A cache coherency protocol employs a cache coherency model to ensure data is correctly maintained throughout the system so that data that is modified is written back to the main memory at an appropriate time to avoid later processing of incorrect data.
[0005] In a multicore system with shared memory, compute elements can access data at a memory address of the main memory, store a copy of the data in one or more levels of cache (e.g., a local cache), modify the copied data, and write the modified data back to the memory address of the main memory. In multicore systems, the cache coherency protocol ensures that all compute elements accessing the same memory location see the most recent value written to it, and, thus, do not operate on an old version of data that has been modified. Data consistency refers to the overall ordering of memory operations across different memory locations to guaranty that all processor circuitry writes in a consistent sequence, even if they are accessing different data items. Thus, coherency corresponds to maintaining the same value for a single data location accessible to multiple cores, while consistency corresponds to the overall ordering of writes to any memory location across the system.

Description

BRIEF DESCRIPTION OF THE DRAWINGS
[0006] FIG. **1** is a block diagram of an example environment in which an example computing system operates to facilitate use of dynamic coherency models for memory objects.
[0007] FIG. **2** is a block diagram of an example implementation of one of the threads of FIG. **1**.
[0008] FIG. **3** is a block diagram of an example implementation of the compute element of FIG. **1**.
[0009] FIG. **4** is a block diagram of an example implementation of the caching agent of FIG. **1**.
[0010] FIGS. **5A-5**E illustrate example implementations of the caching agent of FIG. **1** in different locations of a computing system.
[0011] FIG. **6** is a flowchart representative of example machine readable instructions and/or example operations that may be executed, instantiated, and/or performed by example programmable circuitry to implement the compute element of FIG. **3**.
[0012] FIG. **7** is a flowchart representative of example machine readable instructions and/or

example operations that may be executed, instantiated, and/or performed by example programmable circuitry to implement the caching agent of FIG. **4**.

[0013] FIG. **8** illustrates an example hardware arrangement of an example data center.

[0014] FIG. **9**A illustrates an example arrangement of an example chip assembly of FIG. **8**.

[0015] FIG. **9**B illustrates an example arrangement of an example chip assembly of FIG. **8**, adapted for high-performance computing applications.

[0016] FIG. **10** is a block diagram of an example processing platform including programmable circuitry structured to execute, instantiate, and/or perform the example machine readable instructions and/or perform the example operations of FIGS. **6** and/or **7** to implement one or more of the thread(s), the compute element, or the caching agent of FIGS. **2-4**.

[0017] FIG. **11** is a block diagram of an example implementation of the programmable circuitry of FIG. **10**.

[0018] FIG. **12** is a block diagram of another example implementation of the programmable circuitry of FIG. **10**.

[0019] FIG. **13** is a block diagram of an example software/firmware/instructions distribution platform (e.g., one or more servers) to distribute software, instructions, and/or firmware (e.g., corresponding to the example machine readable instructions of FIGS. **6** and/or **7**) to client devices associated with end users and/or consumers (e.g., for license, sale, and/or use), retailers (e.g., for sale, re-sale, license, and/or sub-license), and/or original equipment manufacturers (OEMs) (e.g., for inclusion in products to be distributed to, for example, retailers and/or to other end users such as direct buy customers).

[0020] In general, the same reference numbers will be used throughout the drawing(s) and accompanying written description to refer to the same or like parts. The figures are not necessarily to scale.

DETAILED DESCRIPTION

[0021] The following introduces examples of computer hardware for facilitate use of dynamic coherency models for memory objects, applicable in programmable architectures such as chiplet-based processors, System-on-chip (SoC) circuitry, System-in-Package (SiP) or System-on-Package (SoP) circuitry, and/or any other implementations of programmable circuitry.

[0022] As used herein, a chiplet refers to any integrated circuit (IC) that has a modular structure to perform one or more specified functionalities and to be combinable with one or more other chiplets on an interposer or other substrate in a package. Examples of chiplets are compute chiplets that include programmable circuitry (e.g., one or more processor circuits, such as one or more cores, etc.) and supporting circuitry (e.g., local memory, etc.) to provide computational functionality (e.g., to execute a host OS, applications, etc.), memory chiplets that include memory accessible to one or more other chiplets, communication chiplets that include communication interfaces (e.g., input/output hubs, networks, etc.) to enable other chiplets to communicate with each other and/or to other devices external to the package, etc. Example multi-tier management architectures provide a flexible management architecture that is multi-tiered to enable management of chiplet-based compute devices that include various combinations of chiplets from various manufacturers. Example implementations of chiplets are further described below in conjunction with FIGS. **8**, **9**A, and **9**B.

[0023] Hardware architecture and programming models perform operations by accessing data from a cache and/or memory, manipulating the accessed data, erasing the accessed data, overwriting the accessed data, etc. For example, a thread may include an instruction to read, modify, and/or write data corresponding to a particular memory address location. When a compute element of a computing system executes the thread, the compute element will access the data (or a copy of the data) corresponding to the memory address location of memory, manipulate the data in a cache or a buffer, and then write the manipulated data back to the memory at the memory address location. A compute element is a component, device, or other circuitry capable of executing and/or

instantiating instructions or issuing execution results. A compute element may be a core, an accelerator, a processor, an accelerator, a central processing unit, a graphical processing unit, etc. Data can be retrieved and/or modified from various levels of the cache/memory system. The cache coherency protocol implements a cache coherency model to maintain correct data throughout the various levels of cache through various operations (e.g., marking data as "dirty," flushing the cache, etc.).

[0024] This process becomes more complicated in a multicore, multichip, multi-chiplet, etc. system where data from a memory address can be accessed by different compute elements. For example, if different compute elements access data from the same memory location and separately manipulate the data in their local cache, the data in the two different caches for the same memory address will be different, leading to coherency issues. Coherency means that data stored in multiple caches are consistent and synchronized. To ensure coherency, a caching agent can utilize one or more coherence protocols, such as modified owned exclusive shared invalid (MOESI), modified shared invalid (MSI), modified shared exclusive invalid (MSEI), modified owned exclusive shared invalid forward (MOESIF), directly-based, token coherence, etc. Coherence protocols ensure that multiple compute elements accessing shared memory always see the most updated (correct) version of data, even if they have copies cached locally, by coordinating cache operations and notifying other compute elements when data is modified (e.g., dirty), and preventing inconsistencies across the computing device. Coherency models may leverage snooping on bus traffic to detect changes and sending invalidation or update messages to other caches storing the data.

[0025] Although sections of memory can be set up for coherency, through the use of coherency models, sections of memory can be set up for no coherency (also referred to as consistency). Consistency (or memory ordering) is based on the accuracy of data in a cache. As used herein, cache refers to one or more levels of a cache hierarchy including the main memory. Programmers develop instructions that ensure consistency without relying on the coherency models. For example, instructions can be written to ensure that two compute elements do not manipulate data in their local cache that correspond to the same memory address location at the same time. In this manner, before a compute element manipulates data for a memory location in its local cache, the data being manipulated in the cache is the same as the data in the main memory (i.e., the data is unchanged).

[0026] In some examples, hardware creates dependencies or performance glass jaws (e.g., ping pong effects) between threads implemented by different compute elements that may access similar parts of data stored in the coherent domain. A glass jaw and/or ping pong occurs when multiple compute elements repeatedly access and modify the same cache line, causing the cache line to bounce back and forth between respective caches. Such glass jaw and/or ping pong events cause corresponding cache line to be updated and invalidated in other compute element caches. However, coherency models to handle glass jaw and/or ping pong events may only be needed for a certain amount of time or to certain memory ranges based on how an application was written. Applications utilize data sets of different natures. In a multi-thread application, some of those data sets may only be used for one single thread or compute element. The same data objects may need or not need coherency across threads depending on the phase of the application.

[0027] In traditional computer systems, sections of memory can be set up for coherency, consistency, and/or a combination of coherency and consistency based on the code being implemented. However, once defined, the coherency/consistency mode of the sections of memory does not change in traditional computer systems. For example, if, for a given program, the memory is set up for consistency, the consistency will not change throughout the program. Although coherency may be needed for only a portion of time, if coherency is needed, then the memory is tagged for a coherence protocol for the entire time. However, utilizing a coherence protocol substantially increases overhead.

[0028] Reduced instruction set computer (RISC)-V architectures includes portions that are

dynamically parametrizable. Examples disclosed herein utilize this dynamic parameterizable functionality of RISC-V to implement dynamic coherency models for memory. For example, using RISC-V, examples disclosed here can create instruction set architecture (ISA) extensions providing a new type of instruction and semantics. An ISA generally refers to rules that describe how a particular form of processor circuitry (such as a compute element) performs operations. For example, an ISA may define the types of machine-readable instructions, the maximum length of a machine-readable instruction, and the format of machine-readable instructions, which are supported by the compute element(s). The compute element(s) can therefore execute any instructions written in a high-level programming language that, when compiled, conforms to instructions in a low-level programmable language that comply with the rules set by the ISA.

[0029] Examples disclosed herein mitigate collateral ping pong effects of coherency by providing a mechanism that allows applications to dynamically tag certain memory objects (e.g., sections of memory) as consistent or coherent for durations of time. In this manner, memory objects can be tagged for coherency when needed and tagged for consistency when coherency is not needed, thereby conserving resources. For example, if an application is written in a manner to avoid coherency issues, the program can be executed in consistent memory, where coherency models are not implemented because the program was written to ensure coherency. However, if the code was not written to ensure coherency, then a coherency model is needed to ensure coherency. In some examples, hardware can detect strange, wrong, or inconsistent behaviors (e.g., too much ping) and notify the software of the behaviors to avoid coherency and/or consistency issues. In such examples, the software may instruct the hardware to look for particular issues in the program and/or to ignore other types of issues in the program. Examples disclosed herein provide an expansion of instructions and operations of a computing system to enable programmers to utilize dynamic and flexible use of memory according to application needed. Additionally, examples disclosed herein could be applied to coherency models between compute elements and other types of devices (e.g., input output (IO) devices connected to CXL), across fabric, to different System-on-chips (SoCs), across chiplets, etc.

[0030] FIG. **1** illustrates an example computing system **100** (e.g., a package, a platform, etc.) to facilitate use of dynamic coherency models for memory objects. The computing system **100** includes an example system-on-chip **102**. The system-on-chip **102** includes an example I/O network card **104**, example threads **106**, **108**, an example compute element **110** (which may execute the threads **106**. **108** in parallel), an example cache **112**, an example caching agent **114**, and an example memory controller **116**. The system-on-chip **102** is connected to example memory channels (e.g., dual inline memory modules (DIMMs) **118***a*, **118***b*, **118***c* which may operate as main memory in the cache hierarchy. Although the computing system **100** of FIG. **1** includes two threads, a single compute element, a single cache, and three memory channels, the computing system **100** may include any number or type of threads, compute elements, caches, and/or memories. For example, although the example computing system **100** of FIG. **1** is shown with two threads, there can be any number of threads.

[0031] The I/O network card **104** of FIG. **1** acts as a bridge between the system-on-chip **102** and/or other devices (e.g., chips, platforms, I/O devices, chiplets, etc.). For example, the I/O network card **104** can send or receive data by facilitating communication with other devices in the system and/or on a network.

[0032] The threads **106**, **108** of FIG. **1** are a sequences of instructions (e.g., corresponding to software and/or firmware) that the compute element **110** executes. The threads may exist at the operating system level or application level and can be managed by a thread scheduler. As further described below, the compute element **110** can execute the threads **106**, while employing one or more dynamic coherency models for memory objects. As further described below in conjunction with FIG. **2**, the threads **106**, **108** may include memory runtime, library application programming interfaces (APIs), and/or operating system page tables for purposes of the dynamic coherency

models and security.

[0033] The compute element **110** of FIG. **1** may be implemented by programmable circuitry such as one or more chiplets, one or more cores, one or more accelerators, one or more microprocessors (e.g., a CPU, a GPU, a VPU and/or an XPU), one or more FPGAs and/or other hardware that executes instructions (e.g., the threads **106**, **108**). The compute element can execute an instruction from the threads **106**, **108** to configure a new memory range to a particular coherency model for one or more threads based on the instruction so that the data from the memory range is handled using the particular coherency model. The compute element **110** may execute the instruction in ring 0 (e.g., by an operating system). Thus, the operating system may ensure that the application is setting up the configurations to resources owned by the application. The instruction may include a memory range associated with the memory mode (e.g., coherency or consistency), a memory type (e.g., a coherence protocol to utilize for the memory range), two IP addresses that correspond to execution of the current memory mode (e.g., after the end address ends, the previous memory mode or a default memory mode is established), a bitmask indicating which cores or hardware threads and/or compute elements are mapped into the memory mode, a bitmask indicating which chiplets and/or IO devices are part of the coherency model, a service level objective with quality of service (e.g., memory bandwidth associated with the memory range, and/or the type of coherence protocols for the memory range. A bitmask is a binary number that uses binary values (e.g., flags to identify the existence of a feature or trait, such as which compute elements or threads are mapped to a memory mode. A bitmask may also be referred to as a bit field, a bit set, a bit array, a binary field, a binary set, a binary array, binary feature identifier, a bit feature identifier, bit flags, binary flags, etc. The compute element **110** of this example tracks the IP address(es) for the memory modes associated to a particular range of instructions. In this manner, the compute element **110** can revert and/or reset the memory mode for the address range to a previous or default coherency protocol and/or model (e.g., consistency, no coherency, etc.) when the end instructions corresponding to the IP address is executed. When the compute element **110** is executing an instruction to access data corresponding to a memory address, the compute element **110** first determines if a copy of the data for the memory address is stored in the cache **112**. If a copy is stored in the cache **112**, the compute element **110** may perform one or more operations to the data based on the instructions. If a copy is not stored in the cache **112** (e.g., corresponding to a miss in the local cache), the compute element **110** instructs the memory controller **116** to access a copy of the data from the memory address line in one or more of the memory channel DIMMs **118***a*, **118***b*, **118***c*. The memory controller then writes the requested data into one or more levels of the cache hierarchy. An example implementation of the compute element **110** is further described below in conjunction with FIG. **3**.

[0034] The example cache **112** of FIG. **3** is implemented by a temporary volatile memory in which the compute element **110** can store, access, and/or manipulate data associated with a particular memory address. The cache may include one or more levels of memory (e.g., L1 cache, L2 cache, L3 cache, etc.). In some examples, the cache **112** stores a copy of data from a particular memory address of the main memory (e.g., one or more of the memory channel DIMMs **118***a*, **118***b*, **118***c*), so that the compute element **110** can quickly and efficiently access and/or manipulate the data. After the compute element **110** is done operating on and/or otherwise using the data, the data in the cache is flushed (e.g., evicted, written, etc.) to the memory address of the corresponding memory channel DIMM(s) **118***a*, **118***b*, **118***c* (e.g., stored in the main memory). In some examples, the cache **112** may store a shared data set that is shared by different compute elements, threads, chiplets, etc. The cache **112** may be any level and/or type of cache (e.g., micro cache, level 1 cache, level 2 cache, level 3 cache, last level cache, etc.).

[0035] The caching agent **114** of FIG. **1** facilitates the use of dynamic coherency models for memory objects by managing memory configurations owning an address range. For example, the caching agent **114** configures a coherency setting (e.g., corresponding to coherency or no coherency) for a memory range by tagging and/or tracking the memory range (e.g., corresponding

to a memory object) as coherent or consistent (e.g., not coherent) and/or for which threads/compute elements the memory object is coherent or consistent. Additionally, when a memory object is tagged for coherency, the caching agent **114** can track and/or monitor the type of coherency model used.

[0036] The caching agent **114** of this example includes a hardware API that allows the compute element **110** to register the memory mode address space. The API registers memory modes into a system address decoder, as further described below. In the case of failure on creation, the caching agent **114** returns the error to the compute element **110**. However, the returning of the error may result in a conflict with another definition with another set of compute elements and/or a conflicting definition with another memory mode range. Accordingly, the system address decoder registers the dynamic memory space to identify potential conflicting memory range definitions (e.g., across different sets of compute elements or between two registrations). Additionally, the caching agent **114** may execute memory requests from the compute element **110** based on the requested memory mode. In the example of FIG. **1**, the caching agent **114** emulates no coherency from the compute element flow perspective to minimize computing impact on the compute element **110**. For example, if an application says that a memory range is no coherent for a first thread running on the compute element **110**, if the first thread requests ownership of a memory line within the memory range, the compute element **110** will not send the request to the caching agent **114**. Rather, the compute element **110** return to an execution unit of the compute element an acknowledgement mimicking the response that the caching agent **114** would have provided to save coherency messages between the compute element and the caching agent **114** In some examples, functionality of the caching agent **114** may be implemented in the compute element **110**. In such examples, compute element **110** may always be aware of the coherency states of each memory object. The caching agent **114** is further described below in conjunction with FIG. **4**.

[0037] The memory controller **116** of FIG. **1** may be a single memory controller or multiple memory controllers (e.g., for different memories of the computing system **100**). Based on instructions from the compute element **110**, the memory controller **116** can read and/or write data to a memory address line of one or more of the memory channel DIMMs **118a**, **118b**, **118c**. For example, for accessing, the memory controller **116** can determine which memory channel DIMMs **118a**, **118b**, **118c** corresponds to a memory address line, read the data from the corresponding DIMM, and store a copy of the read data into the cache **112**. Additionally, for eviction, flushing, writing, etc. the memory controller **116** can determine which memory channel DIMMs **118a**, **118b**, **118c** corresponds to a memory address line, and write data from the cache **112** corresponding to the memory address line into the determined memory channel DIMM **118a**, **118b**, **118c** at the memory address line.

[0038] The memory channel DIMMs **118a**, **118b**, **118c** of FIG. **1** include random access memory. Each memory cell of the memory channel DIMMs **118a**, **118b**, **118c** corresponds to a memory address. Data is stored into a location corresponding to a memory address of the memory channel DIMMs **118a**, **118b**, **118c**. In some examples, the memory channels DIMMs **118a**, **118b**, **118c** may be replaced with one or more different types of memory channels.

[0039] FIG. **2** is a block diagram of an example implementation one or more of the threads **106**, **108** of FIG. **1**. The threads **106**, **108** of FIG. **2** may be instantiated (e.g., creating an instance of, bring into being for any length of time, materialize, implement, etc.) by instructions executed by programmable circuitry. For example, programmable circuitry may be implemented by a Central Processor Unit (CPU) a chiplet, an array of chiplets, a programmable logic device (PLD), a generic array logic (GAL) device, a programmable array logic (PAL) device, a complex programmable logic device (CPLD), a simple programmable logic device (SPLD), a microcontroller (MCU), a programmable system on chip (PSoC), etc. Additionally or alternatively, the thread(s) **106**, **108** of FIG. **2** may be instantiated (e.g., creating an instance of, bring into being for any length of time, materialize, implement, etc.) by (i) an Application Specific Integrated Circuit (ASIC) and/or (ii) a

Field Programmable Gate Array (FPGA) (e.g., another form of programmable circuitry) structured and/or configured in response to execution of second instructions to perform operations corresponding to the first instructions. It should be understood that some or all of the circuitry of FIG. **2** may, thus, be instantiated at the same or different times. Some or all of the circuitry of FIG. **2** may be instantiated, for example, in one or more threads executing concurrently on hardware and/or in series on hardware. Moreover, in some examples, some or all of the circuitry of FIG. **2** may be implemented by microprocessor circuitry executing instructions and/or FPGA circuitry performing operations to implement one or more virtual machines and/or containers. The threads **106**, **108** include example memory runtime and libraries APIs **200** and example operating system page table circuitry **202**.

[0040] The memory runtime and libraries APIs **200** include programming language that provides a mechanism (e.g., via pragmas) to allow an instruction to specify memory ranges or objects that are subject to a particular memory mode and to configure different parameters for the memory range or object. As used herein, a pragma is a directive that provides additional information to the compiler to specify how a compiler should process its input, allowing a programmer to influence specific aspects of the compilation process. A pragma enables specific control over the compilation process beyond what is traditionally available in the programming language itself. The memory runtime and libraries APIs **200** create the address space coherency type instruction, that when executed by a compute element, cause a memory object to temporarily correspond to a particular type of coherent protocol for one or more threads/compute elements.

[0041] The operating system page table **202** of FIG. **2** includes information including the coherency domains that are created when executed and/or implemented, which compute elements and/or threads own the coherency domains, etc. The operating system page table **202** includes such information to avoid security gaps (e.g., unauthorized or unintended use) by ensuring that the software stack requiring changes to a memory range corresponds to (e.g., is owned by) the process and the cores that are mapped to the process are authorized for the process. The operating system may be responsible for allocation of objects and check that the software stack is using the objections without violating other applications. A software stack includes the operating system, a compiler, libraries, and/or an application. The operating system may be operating within the threads **106**, **108** and/or beneath the threads. The operating system may be part of a virtual machine (VM) and/or may be outside of the virtual execution environment (e.g., in the context of a container).

[0042] FIG. **3** is a block diagram of an example implementation of the compute element **110** of FIG. **1** to cause memory ranges to implement dynamic coherency models. The compute element **110** of FIG. **3** may be instantiated (e.g., creating an instance of, bring into being for any length of time, materialize, implement, etc.) by programmable circuitry. For example, programmable circuitry may be implemented by a Central Processor Unit (CPU) a chiplet, an array of chiplets, a programmable logic device (PLD), a generic array logic (GAL) device, a programmable array logic (PAL) device, a complex programmable logic device (CPLD), a simple programmable logic device (SPLD), a microcontroller (MCU), a programmable system on chip (PSoC), etc. Additionally or alternatively, the compute element **110** of FIG. **3** may be instantiated (e.g., creating an instance of, bring into being for any length of time, materialize, implement, etc.) by (i) an Application Specific Integrated Circuit (ASIC) and/or (ii) a Field Programmable Gate Array (FPGA) (e.g., another form of programmable circuitry) structured and/or configured in response to execution of second instructions to perform operations corresponding to the first instructions. It should be understood that some or all of the circuitry of FIG. **3** may, thus, be instantiated at the same or different times. Some or all of the circuitry of FIG. **3** may be instantiated, for example, in one or more threads executing concurrently on hardware and/or in series on hardware. Moreover, in some examples, some or all of the circuitry of FIG. **3** may be implemented by microprocessor circuitry executing instructions and/or FPGA circuitry performing operations to implement one or more virtual machines and/or containers. The compute element **110** includes example interface circuitry **300**,

example memory mode tracking circuitry **302**, and example instruction extension execution circuitry **306**.

[0043] The interface circuitry **300** of FIG. **3** interfaces with the other components of the system-on-chip **102**. For example, the interface circuitry **300** can obtain instructions, can transmit requests and/or receive acknowledgments and/or other data from/to the caching agent **114** and/or from/to the thread(s) **106**, **108**, access and/or manipulate data from/to the cache **112** and/or main memory, transmit instructions to the memory controller **116**, etc. In some examples, the interface circuitry **300** is instantiated by programmable circuitry executing memory access monitoring instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **6**.

[0044] The memory mode tracking circuitry **302** of FIG. **3** tracks the IP address (e.g., the starting memory address for a particular memory range) for the memory modes associated with a particular range of instructions. As further described below, the instruction extension execution circuitry **306** can execute coherency instructions from one or more of the threads **106**, **108**. The coherency instruction may correspond to a set coherency instruction or a reset (e.g., revert) coherency instruction. The memory mode tracking circuitry **302** determines when a set coherency instruction has been obtained for a memory range and monitors the execution of instructions for the reception of an end coherency instruction (also referred to as a reset coherency instruction) for the memory range (which identifies the end of coherency for the memory range). When a set coherency instruction has been executed/obtained, the memory mode tracking circuitry **302** transmits a coherency request to the caching agent **114** to tag the memory range for coherency based on the coherency instruction. After the memory mode tracking circuitry **302** determines that end coherency instruction (e.g., a revert instruction that signifies the end of the coherency mode for a memory range) for a memory range has been obtained/executed, the memory mode tracking circuitry **302** reverts the memory mode for the memory range by transmitting a request to the caching agent **114** that identifies the memory range to revert back to a default coherency protocol and/or model (e.g., no coherency). In some examples, the memory mode tracking circuitry **302** is instantiated by programmable circuitry executing memory access monitoring instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **6**.

[0045] The coherency determination circuitry **304** of FIG. **3** can determine the type of coherency that a given address or address range has. For example, the coherency determination circuitry **304** can communicate, via the interface circuitry **300**, with the caching agent **114** to determine coherency types of address ranges (e.g., by access a stored data in the caching agent **114** that tracks coherency of memory ranges). In this manner, the coherency determination circuitry **304** can prevent access to data in a memory address for a clashing coherency mode (e.g., when the compute element **110** executes a thread that assumes coherency for a memory address that is not tagged for coherency or vice versa).

[0046] The instruction extension execution circuitry **306** of FIG. **3** executes instructions corresponding to memory mode coherency for a memory range (e.g., a memory object). For example, when the thread **106**, **108** includes a coherency instruction, the instruction extension execution circuitry **306** executes the coherency instruction to set or reset coherency for a memory range based on the coherency instruction. The instructions are executed via a system call in the operating system (OS) implemented in the computer elements of the computing system **100**. The OS may check rights for the application. For example, an application may generate a system call to change memory coherency, then the OS performs corresponding checks (e.g., to ensure the application owns the memory range of the system call, and then the OS performs execution of the instruction. An example of different set and/or reset instructions and/or ISAs is shown in the below Table 1

TABLE-US-00001 TABLE 1 Coherency instructions # Instruction Type Instruction 1 Setup a coherent protocol for set_ca_mem @x, offset, a particular memory range coherency_model 2 Setup a coherent protocol for set_ca_mem_cores @x, offset a particular memory range for

coherency_model, core-bitmask a set of cores 3 Setup a coherent protocol for set_ca_mem_chiplet @x, offset a particular memory range for coherency_model, chiplet-bitmask a set of chiplets 4 Setup a coherent protocol for set_ca_mem_ce_io @x, offset a particular memory range for coherency_model, core-bitmask, a set of cores and I/O devices IO-bitmask 5 Setup a coherent protocol for set_ca_mem_ce_io @x, offset a particular memory range for coherency_model, chiplet-bitmask, a set of chiplet and I/O IO-bitmask devices 6 Reset or revert the coherent Reset_ca @x protocol for a particular address range.

[0047] In the above Table 1, @x corresponds to the starting memory address of a memory object or memory range for the memory mode, offset corresponds to the size of the memory object or memory range for the memory mode (e.g., the offset from the starting memory address), and bitmask corresponds to the set of components (e.g., compute elements/threads, chiplets, IO devices) that the memory mode applies to. The instruction extension execution circuitry **306** executes the first coherency instruction to set up a memory range for a particular coherency model for all compute elements and/or threads in the computing system **100**. The instruction extension execution circuitry **306** executes the second coherency instruction to set up a memory range for a particular coherency model for a subgroup of compute elements and/or threads in the computing system **100**. The instruction extension execution circuitry **306** executes the third coherency instruction to set up a memory range for a particular coherency model for a subgroup of chiplets in the computing system **100**. The instruction extension execution circuitry **306** executes the fourth coherency instruction to set up a memory range for a particular coherency model for a subgroup of compute elements and/or IO devices in the computing system **100**. The instruction extension execution circuitry **306** executes the fifth coherency instruction to set up a memory range for a particular coherency model for a subgroup of chiplets and/or IO devices in the computing system **100**. The instruction extension execution circuitry **306** executes the sixth coherency instruction to reset or revert a memory range to a default coherency protocol. In some examples, the default coherency protocol and/or model corresponds to consistency/no coherency. However, the default coherency protocol and/or model may correspond to a different coherency protocol.

[0048] Although the bitmask of the second instruction of Table 1 identifies a set of compute elements to be implemented by a coherency model, the second instruction may identify a set of threads to be implemented by a coherency model. For example, the caching agent **114** could implement a variance where the coherency setting can be per thread set. In such an example, there may be different policies for different threads (e.g., an event for different applications and/or processes). For example, a first process identifier for a first thread is not coherent for memory range [a, b] and a second process identifier for the first thread is coherent for memory range [a, b]. In such examples, the caching agent can store a list of ranges indexed per process identifier and thread identifier. In some examples, the instruction extension execution circuitry **306** is instantiated by programmable circuitry executing memory access monitoring instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **6**.

[0049] In some examples, model specific registers (MSRs) can be used instead or, or in addition to, instructions or ISAs. In such examples, a first MSR (MSR1) can operate as a trigger, to trigger actions and additional MSRs can provide parameters, as shown in the below Table 2.

TABLE-US-00002 TABLE 2 Coherency-based MSRs # Action Type MSRs 1 Setup a coherent protocol MSR1 = op01 for a particular memory range MSR2 = @x MSR3 = offset MSR4 = coherency_model 2 Setup a coherent protocol MSR1 = op02 for a particular memory MSR2 = @x range for a set of cores MSR3 = offset MSR4 = coherency_model MSR5 = core-bitmask 3 Setup a coherent protocol MSR1 = op03 for a particular memory MSR2 = @x range for a set of chiplets MSR3 = offset MSR4 = coherency_model MSR5 = chiplet -bitmask 4 Setup a coherent protocol MSR1 = op04 for a particular memory MSR2 = @x range for a set of cores MSR3 = offset and I/O devices MSR4 = coherency_model MSR5 = IO -bitmask 5 Setup a coherent protocol MSR1 = op05 for a particular memory MSR2 = @x range for a set of MSR3 = offset chiplet and I/O devices

MSR4 = coherency_model MSR5 = chiplet -bitmask MSR6 = IO-bitmask 6 Reset the coherent protocol MSR1 = op06 for a particular address range. MSR2 = @x

[0050] In such examples, the instruction extension execution circuitry **306** can set values in the MSRs to configure a memory range for coherency based on an instruction from a thread. For example, if the instruction extension execution circuitry **306** obtains an instruction to set a coherency protocol to implement a coherency model for a particular memory range for a set of cores, the instruction extension execution circuitry **306** stores a value corresponding to the op02 in the MSR1, sets the starting address line for the coherency model in the MSR2, sets the offset for the end address line corresponding to the memory range in the MSR3, stores a value corresponding to the coherency model in the MSR4, and stores a core bitmask identifying which cores to need to operate the coherency model in the MSR5. In such examples, the caching agent **114** can identify the details of a coherency model for a particular memory range based on the values stored in the MSRs. The MSRs may be located in the compute element **110**, in the caching agent **114** and/or in another location of the system on chip **102**. Some examples can track coherency by mapping the actual configuration and the status registers of memory, cache, NoC components, etc. and control them directly. Such examples can be implemented by logic within the system on chip **102** to configure into the software (e.g., a kernel, a driver, etc.) at low-level and expose a high-level API with commands as expressed in the above Table 2.

[0051] In some examples, the compute element **110** includes means for transmitting, means for tracking, means for determining coherency and means for executing. For example, the means for transmitting may be implemented by the interface circuitry **300**, the means for tracking may be implemented by the memory mode tracking circuitry **302**, the means for determining coherency may be implemented by the coherency determination circuitry **304**, and the means for executing may be implemented by the instruction extension execution circuitry **306**. In some examples, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, and/or the instruction extension execution circuitry **306** may be instantiated by programmable circuitry such as the example programmable circuitry **1012** of FIG. **10**. For instance, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, and/or the instruction extension execution circuitry **306** may be instantiated by the example microprocessor **1100** of FIG. **11** and/or the chiplet of FIGS. **9**A and/or **9**B executing machine executable instructions such as those implemented by at least blocks **602-616** of FIG. **6**. In some examples, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, and/or the instruction extension execution circuitry **306** may be instantiated by hardware logic circuitry, which may be implemented by an ASIC, XPU, or the FPGA circuitry **1200** of FIG. **12** configured and/or structured to perform operations corresponding to the machine readable instructions. Additionally or alternatively, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, and/or the instruction extension execution circuitry **306** may be instantiated by any other combination of hardware, software, and/or firmware. For example, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, and/or the instruction extension execution circuitry **306** may be implemented by at least one or more hardware circuits (e.g., processor circuitry, discrete and/or integrated analog and/or digital circuitry, an FPGA, an ASIC, an XPU, chiplet(s), core(s), a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) configured and/or structured to execute some or all of the machine readable instructions and/or to perform some or all of the operations corresponding to the machine readable instructions without executing software or firmware, but other structures are likewise appropriate.

[0052] FIG. **4** is a block diagram of an example implementation of the caching agent **114** of FIG. **1**. The caching agent **114** of FIG. **4** may be instantiated (e.g., creating an instance of, bring into being for any length of time, materialize, implement, etc.) by programmable circuitry. For example, programmable circuitry may be implemented by a Central Processor Unit (CPU) a chiplet, an array

of chiplets, a programmable logic device (PLD), a generic array logic (GAL) device, a programmable array logic (PAL) device, a complex programmable logic device (CPLD), a simple programmable logic device (SPLD), a microcontroller (MCU), a programmable system on chip (PSoC), etc. Additionally or alternatively, the caching agent **114** of FIG. **4** may be instantiated (e.g., creating an instance of, bring into being for any length of time, materialize, implement, etc.) by (i) an Application Specific Integrated Circuit (ASIC) and/or (ii) a Field Programmable Gate Array (FPGA) (e.g., another form of programmable circuitry) structured and/or configured in response to execution of second instructions to perform operations corresponding to the first instructions. It should be understood that some or all of the circuitry of FIG. **4** may, thus, be instantiated at the same or different times. Some or all of the circuitry of FIG. **4** may be instantiated, for example, in one or more threads executing concurrently on hardware and/or in series on hardware. Moreover, in some examples, some or all of the circuitry of FIG. **4** may be implemented by microprocessor circuitry executing instructions and/or FPGA circuitry performing operations to implement one or more virtual machines and/or containers. The caching agent **114** of FIG. **4** includes example interface circuitry **400**, example system address decoder circuitry **402**, example snoop filtering circuitry **404**, and example coherency logic circuitry **406**.

[0053] The interface circuitry **400** of FIG. **4** obtains coherency requests and/or access requests from the compute element **110**. For example, when the compute element **110** executes an instruction to set a range of memory for coherency, the compute element **110** transmits a request to the interface circuitry **400**. Additionally, when the compute element **110** executes an instruction to access and/or manipulate data stored in a memory address of memory, the compute element **110** transmits an access request to the interface circuitry **400**. Additionally, the interface circuitry **400** can transmit acknowledgments to the compute element **110** for the requests. In some examples, the interface circuitry **400** is instantiated by programmable circuitry executing interface instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **7**.

[0054] The system address decoder circuitry **402** of FIG. **4** registers the memory modes for memory ranges based on coherency requests from the compute element **110**. The system address decoder circuitry **402** can be used to identify potential conflicting memory range definitions (e.g., across different sets of cores or between two registrations). After a coherency request is obtained, the system address decoder circuitry **402** processes the coherency request to determine details related to the coherency so that the address decoder circuitry **402** can configure a coherency setting for the memory range based on the coherency request by tagging or otherwise tracking the coherency details and/or implementing coherency for the memory range based on the coherency details. For example, the coherency request includes one or more of a memory address range, a coherent protocol type, an identification (e.g., a bitmask) of compute elements, threads, chiplets, and/or IO devices. After the system address decoder circuitry **402** obtains a reset request, the system address decoder circuitry **402** untags or otherwise determines that coherency has ended for the corresponding memory range. In some examples, the system address decoder circuitry **402** may include storage (e.g., registers) that stores information for tracking memory ranges that have been tagged for coherency for one or more threads and/or compute elements. In this manner, the compute element **110** can access the data to identify information related to coherency of a memory range, as further described above in conjunction with FIG. **3**. In some examples, the system address decoder circuitry **402** is instantiated by programmable circuitry executing interface instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **7**.

[0055] The snoop filtering circuitry **404** of FIG. **4** performs snooping protocols for data access and/or manipulation requests to data corresponding to a memory range that is tagged for coherency. The snoop filtering circuitry **404** can snoop other computing devices (e.g., other compute elements in the same or a different system) that implement threads that have copies of data corresponding to a memory line and give ownership to a particular thread. A snooping protocol may include monitoring a shared bus to detect when another compute element accesses data that is also cached

in the local cache **112**. In this manner, the snooping filter circuitry **404** can invalidate and/or update the local copy of data based on the monitoring. After snooping, the snoop filtering circuitry **404** can give ownership of the memory line to a particular thread and/or compute element and cause the other compute elements to evict the data from their local caches. In this manner, the thread and/or compute element can manipulate the data and ensure that all the copies of the data across compute devices are the same. In some examples, the snoop filtering circuitry **404** is instantiated by programmable circuitry executing interface instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **7**.

[0056] The coherency logic circuitry **406** of FIG. **4** manages access requests (e.g., to read data, write data, modify data, etc.) from the compute element **110** based on memory modes associated with memory ranges. For example, when the compute element **110** transmits a request to write to a memory address, the coherency logic circuitry **406** interfaces with the system address decoder circuitry **402** to determine if the memory address and/or thread that the compute element **110** executes corresponds to coherence (e.g., a memory range that has been tagged for coherency). If the memory address does not correspond to coherence, the coherency logic circuitry **406** grants the request via an acknowledgement. If the memory address corresponds to a coherency memory mode, the coherency logic circuitry **406** facilitates the approval or denial of the request based on the memory mode for the memory address. For example, if the request is a read only request, the coherency logic circuitry **406** grants the request. If the request is a write and/or manipulate request, the coherency logic circuitry **406** only grants the request if the thread, compute element, IO devices, chiplet, etc. making the request is included in the coherency information (e.g., covered by the bitmask, if there is one). If access is granted, the coherency logic circuitry **406** may instruct the snoop filtering circuitry **404** to perform snoop filtering to ensure coherency for the data across caches. In some examples, coherency logic circuitry **406** is instantiated by programmable circuitry executing interface instructions and/or configured to perform operations such as those represented by the flowchart(s) of FIG. **7**.

[0057] In some examples, the caching agent **114** includes means for means for accessing requests, means for configuring coherency settings, means for snooping and/or means of facilitating a request. For example, the means for accessing may be implemented by the interface circuitry **400**, the means for configuring may be implemented by the system address decoder **402**, the means for snooping may be implemented by the snoop filtering circuitry **404**, and the means for facilitating a request may be implemented by the coherency logic circuitry **406**. In some examples, the means for accessing may include means for transmitting and the means for configuring may include means for determining whether a memory address of an access request corresponds to a memory objected marked for coherency. In some examples, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, and/or the coherency logic circuitry **406** may be instantiated by programmable circuitry such as the example programmable circuitry **1012** of FIG. **10**. For instance, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, and/or the coherency logic circuitry **406** may be instantiated by the example microprocessor **1100** of FIG. **11** and/or the chiplet of FIGS. **9**A and/or **9**B executing machine executable instructions such as those implemented by at least blocks **702-720** of FIG. **7**. In some examples, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, and/or the coherency logic circuitry **406** may be instantiated by hardware logic circuitry, which may be implemented by an ASIC, XPU, chiplet(s), core(s), or the FPGA circuitry **1200** of FIG. **12** configured and/or structured to perform operations corresponding to the machine readable instructions. Additionally or alternatively, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, and/or the coherency logic circuitry **406** may be instantiated by any other combination of hardware, software, and/or firmware. For example, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, and/or the coherency logic circuitry **406** may be implemented by at least one or more

hardware circuits (e.g., processor circuitry, discrete and/or integrated analog and/or digital circuitry, an FPGA, an ASIC, an XPU, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) configured and/or structured to execute some or all of the machine readable instructions and/or to perform some or all of the operations corresponding to the machine readable instructions without executing software or firmware, but other structures are likewise appropriate.

[0058] FIGS. **5A-5E** illustrate alternative implementations of the caching agent **114** of FIG. **2**. FIG. **5A** includes the system on chip **102** and the caching agent **114** of FIG. **1**. FIG. **5B** includes an example platform **500**, an example package **502**, and the caching agent **114** of FIG. **1**. FIG. **5C** includes an example package **504**, example compute chiplets **506**, an example IO hub **508**, and the caching agent **114** of FIG. **1**. FIG. **5D** includes an example package **510**, example compute chiplets **512**, an example IO hub **514**, and the caching agent **114** of FIG. **1**. FIG. **5E** includes example packages **518**, **524**, example chiplets **520**, **526**, example IO hubs **522**, **528**, an example caching agent **530**, and the example caching agent **114** of FIG. **1**.

[0059] Each configuration has a potential location of the caching agent **114** for different operations. For example, in FIG. **5A**, the caching agent **114** is implemented in the system on chip **102**, as described above in conjunction with FIG. **1**. In this manner, the caching agent **112** can facilitate dynamic memory modes for compute elements inside the system on chip **102**. In FIG. **5B**, the caching agent **114** is implemented on the platform **500** outside of the package **502** (which may include chiplets, systems-on-chips, IO hubs, etc.). In this manner, the caching agent **112** can facilitate dynamic memory modes for compute elements inside the package **502** and/or between the system on chip of FIG. **5A** and external sources (e.g., IO devices). In FIG. **5C**, the caching agent **114** is implemented within the package **504**. In this manner, the caching agent **114** can implement memory modes across the chiplets **506**. In some examples, such as FIG. **5D**, the caching agent **114** is implemented in the IO hub to implement memory modes across the chiplets **512**. However, the caching agent **114** may additionally or alternatively be implemented in one or more of the chiplets **506**, **512**. In FIG. **5E**, the caching agent **114** of the package **518** communicates with the caching agent **530** of the package **524** using unified payments interface (UPI) or compute express link (CXL) capabilities to implement memory modes for coherency across different chips.

[0060] While an example manner of implementing one or more of the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIG. **1** is illustrated in FIGS. **2**, **3** and/or **4** one or more of the elements, processes, and/or devices illustrated in FIGS. **2**, **3**, and/or **4** may be combined, divided, re-arranged, omitted, eliminated, and/or implemented in any other way. Further, the memory runtime and libraries APIs **2100**, the operating system page table **202**, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, the instruction extension execution circuitry **306**, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, the coherency logic circuitry **406**, and/or, more generally, the example the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4**, may be implemented by hardware alone or by hardware in combination with software and/or firmware. Thus, for example, any of the memory runtime and libraries APIs **2100**, the operating system page table **202**, the interface circuitry **300**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, the instruction extension execution circuitry **306**, the interface circuitry **400**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, the coherency logic circuitry **406**, and/or, more generally, the example the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4**, could be implemented by programmable circuitry, such as one or more chiplets, one or more processor cores, processor circuitry, analog circuit(s), digital circuit(s), logic circuit(s), programmable processor(s), programmable microcontroller(s), graphics processing unit(s) (GPU(s)), digital signal processor(s) (DSP(s)), ASIC(s), programmable logic device(s) (PLD(s)) vision processing units(s) (VPUs), and/or field programmable logic device(s) (FPLD(s)) such as FPGAs in combination with machine readable instructions (e.g., firmware or software). Further

still, the example the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4** may include one or more elements, processes, and/or devices in addition to, or instead of, those illustrated in FIGS. **2**, **3**, and/or **4**, and/or may include more than one of any or all of the illustrated elements, processes and devices.

[0061] Flowchart(s) representative of example machine readable instructions, which may be executed by programmable circuitry to implement and/or instantiate the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4** and/or representative of example operations which may be performed by programmable circuitry to implement and/or instantiate the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4**, are shown in FIGS. **6** and/or **7**. The machine readable instructions may be one or more executable programs or portion(s) of one or more executable programs for execution by programmable circuitry such as the programmable circuitry **1012** shown in the example processor platform **1000** discussed below in connection with FIG. **10** and/or may be one or more function(s) or portion(s) of functions to be performed by the example programmable circuitry (e.g., an FPGA) discussed below in connection with FIGS. **11** and/or **12**. In some examples, the machine readable instructions cause an operation, a task, etc., to be carried out and/or performed in an automated manner in the real world. As used herein, "automated" means without human involvement.

[0062] The program may be embodied in instructions (e.g., software and/or firmware) stored on one or more non-transitory computer readable and/or machine readable storage medium such as cache memory, a magnetic-storage device or disk (e.g., a floppy disk, a Hard Disk Drive (HDD), etc.), an optical-storage device or disk (e.g., a Blu-ray disk, a Compact Disk (CD), a Digital Versatile Disk (DVD), etc.), a Redundant Array of Independent Disks (RAID), a register, ROM, a solid-state drive (SSD), SSD memory, non-volatile memory (e.g., electrically erasable programmable read-only memory (EEPROM), flash memory, etc.), volatile memory (e.g., Random Access Memory (RAM) of any type, etc.), and/or any other storage device or storage disk. The instructions of the non-transitory computer readable and/or machine readable medium may program and/or be executed by programmable circuitry located in one or more hardware devices, but the entire program and/or parts thereof could alternatively be executed and/or instantiated by one or more hardware devices other than the programmable circuitry and/or embodied in dedicated hardware. The machine readable instructions may be distributed across multiple hardware devices and/or executed by two or more hardware devices (e.g., a server and a client hardware device). For example, the client hardware device may be implemented by an endpoint client hardware device (e.g., a hardware device associated with a human and/or machine user) or an intermediate client hardware device gateway (e.g., a radio access network (RAN)) that may facilitate communication between a server and an endpoint client hardware device. Similarly, the non-transitory computer readable storage medium may include one or more mediums. Further, although the example program is described with reference to the flowchart(s) illustrated in FIGS. **6** and/or **7**, many other methods of implementing the example the example the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4** may alternatively be used. For example, the order of execution of the blocks of the flowchart(s) may be changed, and/or some of the blocks described may be changed, eliminated, or combined. Additionally or alternatively, any or all of the blocks of the flow chart may be implemented by one or more hardware circuits (e.g., processor circuitry, chiplet(s), discrete and/or integrated analog and/or digital circuitry, an FPGA, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware. The programmable circuitry may be distributed in different network locations and/or local to one or more hardware devices (e.g., a single-core processor (e.g., a single core CPU), a multi-core processor (e.g., a multi-core CPU, an XPU, a chiplet and/or an array of chiplets, etc.)). As used herein, programmable circuitry includes any type(s) of circuit that may be programmed to perform a desired function such as, for example, a CPU, a core, a chiplet, an array of chiplets, a GPU, a VPU, and/or an FPGA. The programmable

circuitry may include one or more CPUs, one or more cores, one or more chiplets, one or more GPUs, one or more VPUs, and/or one or more FPGAs located in the same package (e.g., the same integrated circuit (IC) package or in two or more separate housings), one or more one or more CPUs, one or more cores, one or more chiplets, one or more GPUs, one or more VPUs, and/or one or more FPGAs in a single machine, multiple CPUs, cores, chiplets, GPUs, VPUs, and/or FPGAs distributed across multiple servers of a server rack, and/or multiple CPUs, cores, chiplets, GPUs, VPUs, and/or FPGAs distributed across one or more server racks. Additionally or alternatively, programmable circuitry may include a programmable logic device (PLD), a generic array logic (GAL) device, a programmable array logic (PAL) device, a complex programmable logic device (CPLD), a simple programmable logic device (SPLD), a microcontroller (MCU), a programmable system on chip (PSoC), etc., and/or any combination(s) thereof in any of the contexts explained above.

[0063] The machine readable instructions described herein may be stored in one or more of a compressed format, an encrypted format, a fragmented format, a compiled format, an executable format, a packaged format, etc. Machine readable instructions as described herein may be stored as data (e.g., computer-readable data, machine-readable data, one or more bits (e.g., one or more computer-readable bits, one or more machine-readable bits, etc.), a bitstream (e.g., a computer-readable bitstream, a machine-readable bitstream, etc.), etc.) or a data structure (e.g., as portion(s) of instructions, code, representations of code, etc.) that may be utilized to create, manufacture, and/or produce machine executable instructions. For example, the machine readable instructions may be fragmented and stored on one or more storage devices, disks and/or computing devices (e.g., servers) located at the same or different locations of a network or collection of networks (e.g., in the cloud, in edge devices, etc.). The machine readable instructions may require one or more of installation, modification, adaptation, updating, combining, supplementing, configuring, decryption, decompression, unpacking, distribution, reassignment, compilation, etc., in order to make them directly readable, interpretable, and/or executable by a computing device and/or other machine. For example, the machine readable instructions may be stored in multiple parts, which are individually compressed, encrypted, and/or stored on separate computing devices, wherein the parts when decrypted, decompressed, and/or combined form a set of computer-executable and/or machine executable instructions that implement one or more functions and/or operations that may together form a program such as that described herein.

[0064] In another example, the machine readable instructions may be stored in a state in which they may be read by programmable circuitry, but require addition of a library (e.g., a dynamic link library (DLL)), a software development kit (SDK), an application programming interface (API), etc., in order to execute the machine-readable instructions on a particular computing device or other device. In another example, the machine readable instructions may need to be configured (e.g., settings stored, data input, network addresses recorded, etc.) before the machine readable instructions and/or the corresponding program(s) can be executed in whole or in part. Thus, machine readable, computer readable and/or machine readable media, as used herein, may include instructions and/or program(s) regardless of the particular format or state of the machine readable instructions and/or program(s).

[0065] The machine readable instructions described herein can be represented by any past, present, or future instruction language, scripting language, programming language, etc. For example, the machine readable instructions may be represented using any of the following languages: C, C++, Java, C-Sharp, Perl, Python, JavaScript, HyperText Markup Language (HTML), Structured Query Language (SQL), Swift, etc.

[0066] As mentioned above, the example operations of FIGS. **6** and/or **7** may be implemented using executable instructions (e.g., computer readable and/or machine readable instructions) stored on one or more non-transitory computer readable and/or machine readable media. As used herein, the terms non-transitory computer readable medium, non-transitory computer readable storage

medium, non-transitory machine readable medium, and/or non-transitory machine readable storage medium are expressly defined to include any type of computer readable storage device and/or storage disk and to exclude propagating signals and to exclude transmission media. Examples of such non-transitory computer readable medium, non-transitory computer readable storage medium, non-transitory machine readable medium, and/or non-transitory machine readable storage medium include optical storage devices, magnetic storage devices, an HDD, a flash memory, a read-only memory (ROM), a CD, a DVD, a cache, a RAM of any type, a register, and/or any other storage device or storage disk in which information is stored for any duration (e.g., for extended time periods, permanently, for brief instances, for temporarily buffering, and/or for caching of the information). As used herein, the terms "non-transitory computer readable storage device" and "non-transitory machine readable storage device" are defined to include any physical (mechanical, magnetic and/or electrical) hardware to retain information for a time period, but to exclude propagating signals and to exclude transmission media. Examples of non-transitory computer readable storage devices and/or non-transitory machine readable storage devices include random access memory of any type, read only memory of any type, solid state memory, flash memory, optical discs, magnetic disks, disk drives, and/or redundant array of independent disks (RAID) systems. As used herein, the term "device" refers to physical structure such as mechanical and/or electrical equipment, hardware, and/or circuitry that may or may not be configured by computer readable instructions, machine readable instructions, etc., and/or manufactured to execute computer-readable instructions, machine-readable instructions, etc.

[0067] FIG. **6** is a flowchart representative of example machine readable instructions and/or example operations **600** that may be executed, instantiated, and/or performed by programmable circuitry to facilitate temporary coherency for a memory range using the compute element **110** of FIG. **3**. The example machine-readable instructions and/or the example operations **600** of FIG. **6** begin at block **602**, at which the memory mode tracking circuitry **302** determines if a coherency instruction has been obtained via the interface circuitry **300** and/or executed by the instruction extension execution circuitry **306**. A coherency instruction is an instruction included in a thread that the instruction extension execution circuitry **306** can execute to tag a section of memory (e.g., a memory object or memory range) with a coherency model for one or more threads, compute elements, chiplet, and/or IO devices. The coherency instruction may correspond to the first five instructions included in the above-Table 1 and/or the values stored in the MSRs for the first five action types in the above Table 2.

[0068] If the memory mode tracking circuitry **302** determines that a coherency instruction has not been obtained/executed (block **602**: NO), control continues to block **610**. If the memory mode tracking circuitry **32** determines that a coherency instruction has been obtained/executed (block **602**: YES), the memory mode tracking circuitry **302** uses the interface circuitry **300** to transmit a set coherency request to the caching agent **114** to configure a coherency setting for a memory object based on the coherency instruction (block **604**). For example, the memory mode tracking circuitry **302** can generate a coherence request identifying the memory range, the type of coherent protocol used, the threads, compute elements, chiplets, and/or IO devices that are mapped to the coherency memory mode, a service level objective, etc. The interface circuitry **300** transmits the coherency request to the caching agent **114**, which tags and/or tracks the coherency of memory objects and facilitates operation based on the tagged and/or tracked coherency of the memory objects.

[0069] At block **606**, the memory mode tracking circuitry **302** determines if a revert coherency (also referred to as a reset coherence or end coherence) instruction has been objected by the interface circuitry **300** and/or executed by the instruction extension execution circuitry **306**. The revert coherency is an instruction that reverts the temporary coherency model for a memory object back to a default coherency (e.g., no coherency) and corresponds to the sixth instruction of the above-Table 1 and/or the MSR values of the sixth action type of the above-Table 2. If the memory

mode tracking circuitry **302** determines that a revert coherency instruction has not been obtained/executed (block **606**: NO), control continues to block **610**. If the memory mode tracking circuitry **302** determines that a revert coherency instruction has been obtained/executed (block **606**: YES), the memory mode tracking circuitry **302** uses the interface circuitry **300** to transmit a revert coherency request to the caching agent **114** to revert the coherency setting for a memory object to a default coherency protocol and/or model corresponding to the memory address range identified in the revert coherency instruction (block **608**).

[0070] At block **610**, the instruction extension execution circuitry **306** determines if an access instruction has been obtained. An access instruction is an instruction to access data corresponding to a memory address or memory line. The access instruction may correspond to a read instruction, a write instruction, a manipulate instruction (e.g., add, multiply, modify, etc.), etc. The execution may be part of the instructions of an application or thread. If the instruction extension execution circuitry **306** determines that the access instruction has not been obtained (block **610**: NO), control returns to block **602**. If the instruction extension execution circuitry **306** determines that the access instruction has been obtained (block **612**: YES), the interface circuitry **300** transmits an access request corresponding to the instruction to the caching agent **114** (block **612**). As further described below, the caching agent **114** can grant or deny the access instruction based on the memory mode for the memory address of the access instruction.

[0071] At block **614**, the instruction extension execution circuitry **204** determines if an acknowledgement granting permission has been obtained from the caching agent **114** via the interface circuitry **300**. If the instruction extension execution circuitry **204** determines that an acknowledgment granting permission has not been received (block **614**: NO), control returns to block **602**. In some examples, the instruction extension execution circuitry **306** can discard the instruction, retry the instruction, and/or save the instruction to retry at a different point in time. If the instruction extension execution circuitry **204** determines that an acknowledgment granting permission has been received (block **614**: YES), the instruction extension execution circuitry **306** performs the access operation (block **616**).

[0072] FIG. **7** is a flowchart representative of example machine readable instructions and/or example operations **700** that may be executed, instantiated, and/or performed by programmable circuitry to facilitate temporary coherency for a memory range using the caching agent **114** of FIG. **4**. The example machine-readable instructions and/or the example operations **700** of FIG. **7** begin at block **702**, at which the system address decoder circuitry **402** determines if a coherency request has been obtained (e.g., accessed, received, etc.) via the interface circuitry **400**. As described above, a coherency request is generated by the compute element **110** when executing a coherency instruction that identifies memory mode configurations for a memory address range. The coherency request may be a set coherency request or a revert coherency request. The set coherency request is a request to configure coherency for a memory object. revert coherency request is a request to remove coherency from a previously set memory object.

[0073] If the system address decoder circuitry **402** determines that a coherency request has not been obtained (block **702**: NO), control continues to block **706**. If the system address decoder circuitry **402** determines that a coherency request has been obtained (block **702**: YES), the system address decoder circuitry **402** applies a coherency model by configuring the coherency setting for the memory range based on the coherency request (block **704**). For example, the system address decoder circuitry **402** can configure the coherency setting for the memory range by tagging or otherwise tracking the coherency details and/or implementing coherency for the memory range based on the coherency details, as further described above in conjunction with FIG. **4**. The system address decoder circuitry **402** identifies the memory range in the coherency request based on the information included in the coherency request. The coherency request may identify the memory range, a coherent protocol type for implementing a coherency model, the coherency model, and/or a bitmask identifying which threads, compute elements, IO devices, and/or chiplets are to utilize

the coherent protocol. The coherency request may be a set coherency request or a revert coherency request. The system address decoder circuitry **402** ensures that coherency is implemented for the memory range identified in the coherency request based on the details of the coherency request until the coherency for the memory range is reset. In some examples, for a set coherency request, the system address decoder circuitry **402** may store an entry corresponding to the coherency request and/or tags a memory object corresponding to a memory address range identified in the coherency request, based on the coherency information included in the coherency request. Alternatively, the system address decoder circuitry **402** may store the coherency request and lookup the coherency request information in response to an access request to a memory address included in the memory range of the coherency request. The system address decoder circuitry **402** may tag and/or track the memory range based on the information corresponding to the coherency request, as opposed to creating an entry. In some examples, the coherency request is a request to end and/or reset coherency for a previously tagged memory range. In such examples, the system address decoder circuitry **402** removes the entry for the memory range, untags the memory range for coherency, and/or stops implementing coherency for the memory range.

[0074] At block **706**, the coherency logic circuitry **406** determines if the interface circuitry **400** has obtained an access request from the compute element **110**. As described above, the access request corresponds to a read operation, a write operation, or a manipulation operation to data at a memory address based on an instruction being executed by the compute element **110**. If the coherency logic circuitry **406** determines that an access request has not been obtained (block **706**: NO), control returns to block **702**. If the coherency logic circuitry **406** determines that an access request has been obtained (block **706**: YES), the system address decoder circuitry **402** determines whether the access request corresponds to a coherency-based memory object (block **708**). A coherency-based memory object is a memory object corresponding to a range of memory addresses that has been tagged or being tracked as corresponding to a coherency model. Because the system address decoder circuitry **402** tracks the coherency-based memory objects based on obtained coherency requests, the system address decoder circuitry **402** can determine if the access request is for a memory address that is within a memory range or memory objected tagged as coherent.

[0075] If the system address decoder circuitry **402** determines that the access request corresponds to a memory object tagged as coherent (block **708**: YES), control continues to block **712**. If the system address decoder circuitry **402** determines that the access request does not correspond to a memory object tagged as coherent (block **708**: NO), the coherency logic circuitry **406** instructs the interface circuitry **400** to transmit an acknowledgement to the compute element **110** granting access to the request (block **710**). In this manner, the compute element **110** can execute the access instruction. At block **712**, the coherency logic circuitry **406** determines if the coherency characteristic for the memory object that includes the memory address identified in the access request corresponds to coherency logic for the compute element that provided the access request. For example, if the memory object corresponds to a bitmask that identifies coherency for particular compute elements, IO devices, chiplets, threads, etc. the coherency logic circuitry **406** determines whether the compute element and/or thread that transmitted the access request is identified by the bitmask as adhering to coherency.

[0076] If the coherency logic circuitry **406** determines that the coherency characteristics correspond to coherency for the compute element and/or thread that provided the access request (block **712**: YES), control continues to block **720**. If the coherency logic circuitry **406** determines that the coherency characteristics do not correspond to coherency for the compute element and/or thread that provided the access request (block **712**: NO), the coherency logic circuitry **406** determines if the access request corresponding to modifying data of the memory object (e.g., such as a write operation, a manipulate operation, a read modify write operation, etc.) (block **714**). If the coherency logic circuitry **406** determines that the access request does not correspond to modifying the memory object (e.g., the access request corresponds to a read operation) (block **714**: NO), the

coherency logic circuitry **406** instructs the interface circuitry **400** to transmit an acknowledgement to the compute element **110** granting the access request (block **716**) and control returns to block **702**. In this manner, the compute element **110** can proceed with the read operation to the memory address.

[0077] If the coherency logic circuitry **406** determines that the access request corresponds to modifying the memory object (block **714**: YES), the coherency logic circuitry **406** instructs the interface circuitry **400** to transmit an acknowledgement to the compute element **110** denying the access request (block **718**) and control returns to block **702**. If the coherency logic circuitry **406** determines that the coherency characteristics corresponds to coherency for the compute element and/or thread that provided the access request (block **712**: YES), the coherency logic circuitry **406** and/or the snoop filtering circuitry **404** facilitate the access request according to the coherency model for the coherency-based object (block **720**) and control returns to block **702**. As described above, the coherency model may include snooping, causing data to be flushed from all other caches for other compute elements, promoting the thread and/or compute element to owner of the memory address for a modification of the data, etc.

[0078] FIGS. **8**, **9**A, **9**B, and **10** include example computing architectures in which any of the techniques and configurations above may be implemented.

[0079] FIG. **8** illustrates an example hardware arrangement of an example data center **800** used to provide multiple examples or instances of a computing system (e.g., the programmable circuitry platform **1000**, described below), with each example of the computing system identified as a respective platform (e.g., the platform **830**, described below). The data center **800** includes example data center infrastructure **801**, an example data center network fabric **802**, and an example power distribution unit **803** to support multiple racks of compute platforms, with a single instance of an example rack **810** depicted. The data center infrastructure **801** may provide physical components that host the compute platform hardware, storage components, and/or networking equipment. The data center network fabric **802** may include switches and/or networking components to support data flows among various compute platforms and storage devices throughout the data center. The power distribution unit **803** may include components to distribute and/or control power among the various compute platforms, networking, and storage devices.

[0080] The rack **810** of FIG. **8** includes, but is not limited to, example cooling infrastructure **811**, an example network interface **812**, and/or other related physical components to support discrete instances of multiple chassis. The rack **810** provides power, connectivity, and/or cooling to each of the multiple chassis in a single rack, with a single instance of a chassis **820** in the example of FIG. **8**. The chassis **820** includes, but is not limited to, example cooling infrastructure **821**, an example chassis network fabric **822**, and an example power supply **823**, which provides cooling, network connectivity, and/or power to multiple platforms within the chassis. Although a single instance of an example platform **830** is illustrated in FIG. **8**, in some examples, a common data center rack configuration may include dozens of chassis, with each chassis to support a number of platforms depending on the physical size of the platform hardware and/or supporting equipment.

[0081] The platform **830** of FIG. **8** may be referred to as a server or node, depending on the use case for the platform **830** and the data center **800**. The platform **830** includes but is not limited to examples of a discrete computing system hosted on a single board. In FIG. **8**, the platform **830** is illustrated as hosting a first example chip assembly **840**A and a second example chip assembly **840**B on a first board provided by a printed circuitry board (PCB) or other platform board, shown as an example PCB **831**. In some examples, the platform **830** may include only one chip package, whereas the PCB **831** includes interconnection of multiple chip assemblies via an interface (e.g., a peripheral component interconnect express (PCIe) interface). Additional chip packages and components may also be hosted on the PCB **831**.

[0082] Some examples of the chip assembly **840**A, **840**B of FIG. **8** may be termed as a System-on-Chip (SoC) package, as modular chiplets that perform different functions are integrated into a

single package-even though this chip package is composed of multiple dies unlike a traditional SoC design that uses a single die. Other examples of the chip assembly **840**A, **840**B may include a System-on-Package (SoP), System-in-a-Package (SiP), or other single chip packages. Various combinations of 2 dimension (D), 2.5D, and/or 3D packaging technologies may be used to manufacture and/or assemble the chip package and its underlying structure. Additionally, different manufacturing processes may be used to provide chiplets and components from different process nodes (e.g., semiconductor fabrication systems).

[0083] The first chip assembly **840**A and the second chip assembly **840**B of FIG. **8** are packages that include multiple chiplets and/or dies for respective functions, such as separate chiplets for processing (e.g., central processing unit (CPU) or graphical processing unit (GPU) chiplets), memory (e.g., cache or high-bandwidth memory chiplets), input/output (I/O) (e.g., I/O chiplets), acceleration (e.g., artificial intelligence (AI)/machine learning (ML) acceleration chiplets), signal processing (e.g., audio or video processing chiplets), etc. The close-up of chip assembly **840**A of FIG. **8** includes a I/O Hub chiplet **841**, chiplets **842**, and a power supply **843**. These components may be hosted on an interposer that is designed to connect multiple dies and/or components within a single semiconductor package (e.g., chip package). In some examples, the chiplets **842** may be manufactured and/or sourced separately and later assembled into the chip package to create the chip assembly **840**A. Various connections may be provided among the chiplets **842**, such as with the use of Universal Chiplet Interconnect Express (UCIe) interfaces and communications, and/or between chiplets and on-chip memory (e.g., high-bandwidth memory (HBM)) using HBM3 (JEDEC), Universal Memory Interface (UMI), or other memory interfaces.

[0084] FIG. **9**A illustrates an example arrangement of an example chip assembly **940**A (e.g., a multi-processing core example of the first chip assembly **840**A or the second chip assembly **840**B of FIG. **8**), with expanded views of the chiplets and processing units included herein. In FIG. **9**A the chip assembly **940**A, which may constitute a SoC, SoP, SiP, and/or other type of chip package, includes chiplets such as an example chiplet **910**A, an example chiplet **910**B, etc. and associated on-package memory (e.g., high-speed memory) such as 3D-stacked, High Bandwidth Memory (HBM) instances (shown as an example HBM **920**A, an example HBM **920**B, interfaces (e.g., UCle interfaces) shown as an example UCle **921**A, an example UCle **921**B, and an example I/O hub **930** (e.g., which may be implemented by a I/O chiplet). Other hardware elements of a chip package are not included for simplicity. Although the examples disclosed herein are described in conjunction with UCLe interfaces, one or more of the interfaces may be device-to-device (Dev2Dev) interfaces (e.g., CXLI, peripheral component interconnect express (PCIE)), die to die (D2D) interfaces (e.g., NVLINK), chiplet to chiplet (Ch2Ch) interfaces (e.g., universal chiplet interconnected express (UCIe)), core to core (C2C) interfaces (e.g., using coherency models), etc.

[0085] The chiplets **910**A, **910**B of FIG. **9**A include multiple processing units and the example processing units **900**A, **900**B, **900**C, **900**D include one or multiple cores, respectively. For example, the chiplet **910**A of FIG. **9**A includes four processing units (the processing units **900**A, **900**B, **900**C, **900**D) and an example Level 3 (L3) cache **904**. The processing units **900**A, **900**B, **900**C, **900**D may include one or multiple processing cores, one or multiple caches, other processing units and/or passive and/or active elements. For example, processing unit **900**A includes two cores (an example core **901**A and an example core **901**B), vector processing unit **902**, and an example level 2 (L2) cache **903**. Accordingly, a single-core processing unit can provide four cores per chiplet and eight total cores in a two-chiplet chip assembly, whereas a dual-core processing unit can provide eight cores per chiplet and sixteen total cores in a two-chiplet chip assembly. However, examples disclosed herein may correspond to other permutations.

[0086] FIG. **9**B is an example arrangement of an example chip assembly **940**B (e.g., a multi-chiplet high-performance computing (HPC) example of chip assembly **840**A, **840**B), adapted for HPC applications (e.g., parallel processing operations involving thousands, millions, or more of processors and/or cores operating simultaneously). The example chip assembly **940**B illustrates

placement as a SiP, SoC, and/or other package onto a platform board (e.g., the PCB **831** of FIG. **8**). The platform board may be in a data center (e.g., the data center **800** of FIG. **8**) or in a standalone deployment setting (e.g., in a standalone computer system, mobile computing device, autonomous device, etc.).

[0087] The chip assembly **940**B of FIG. **9**B is composed of multiple chiplets, shown with four chiplets, including example chiplets **910**C, **910**D, **910**E, **910**F. The chiplets **910**C, **910**D, **910**E, **910**F include multiple processing units, such as thirty-two processing units with a corresponding level 3 (L3) cache for each processing unit. The processing units may include one or multiple cores, such as an example single-core processing unit **900**E shown as part of the chiplet **910**C. The chip assembly **940**B also includes corresponding memory resources, such as HBM elements corresponding to respective banks of processing units (e.g., HBM **920**B and HBM **920**C corresponding respective sets of processing units of chiplet **910**C), UCle interfaces, and/or an IO Hub.

[0088] The chip assembly and related products or devices described herein may be configured in a variety of computing system examples. Such examples include non-transitory machine-readable media storing machine-readable instructions and one or more processors coupled to the memory, such that executing the machine-readable instructions configure one or more of the processors and/or implementing hardware (e.g., the processing unit **900**, the chiplet **910**, the chip **840**, and/or the platform **830** of FIGS. **8**, **9**A, and/or **9**B) to perform operations described above for electronic systems or devices (e.g., to facilitate the use of dynamic memory modes, etc.). It should be further understood that software, including one or more machine readable instructions, that facilitate processing and operations as described above may be distributed, installed, or otherwise provided to networked devices (e.g., servers or cloud computing systems). Alternatively, in some examples, the software may be obtained and loaded (or, re-loaded/upgraded) from one or more servers and/or cloud computing systems, such as software stored on a server for distribution over the Internet, for example.

[0089] FIG. **10** is a block diagram of an example programmable circuitry platform **1000** structured to execute and/or instantiate the example machine-readable instructions and/or the example operations of FIGS. **6** and/or **7** to implement the example the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4**. The programmable circuitry platform **1000** can be, for example, a server, a personal computer, a workstation, a self-learning machine (e.g., a neural network), a mobile device (e.g., a cell phone, a smart phone, a tablet such as an iPad™), a personal digital assistant (PDA), an Internet appliance, a gaming console, or any other type of computing and/or electronic device.

[0090] The programmable circuitry platform **1000** of the illustrated example includes programmable circuitry **1012**. The programmable circuitry **1012** of the illustrated example is hardware. For example, the programmable circuitry **1012** can be implemented by one or more integrated circuits, logic circuits, FPGAs, microprocessors, CPUs, GPUs, DSPs, and/or microcontrollers from any desired family or manufacturer. In some examples, the programmable circuitry **1012** can be implemented by reduced instruction set computer (RISC)-V architecture and/or a chiplet (e.g., the chiplet assemblies **840**A, **840**B, **940**A, **940**B of FIGS. **9**, **10**A and/or **10**B). The programmable circuitry **1012** may be implemented by one or more semiconductor based (e.g., silicon based) devices. In this example, the programmable circuitry **1012** implements the memory runtime and libraries APIs, the operating system page table **202**, the memory mode tracking circuitry **302**, the coherency determination circuitry **304**, the instruction extension execution circuitry **306**, the system address decoder circuitry **402**, the snoop filtering circuitry **404**, and the coherency logic circuitry **406** of FIGS. **2-4**.

[0091] In some examples, the hardware of the circuitry may include variably connected physical components (e.g., execution units, transistors, simple circuits, etc.) including a machine-readable medium physically modified (e.g., magnetically, electrically, moveable placement of invariant

massed particles, etc.) to encode instructions of the specific operation. In connecting the physical components, the underlying electrical properties of a hardware constituent are changed, for example, from an insulator to a conductor or vice versa. The instructions enable embedded hardware (e.g., the execution units or a loading mechanism) to create members of the circuitry in hardware via the variable connections to carry out portions of the specific operation when in operation. Accordingly, the machine-readable medium elements can be part of the circuitry or communicatively coupled to the other components of the circuitry when the device is operating. Also, in some examples, any of the physical components may be used in more than one member of more than one circuitry. For example, under operation, execution units may be used in a first circuit of first circuitry at one point in time and reused by a second circuit in the first circuitry, or by a third circuit in a second circuitry at a different time.

[0092] The programmable circuitry **1012** of the illustrated example includes a local memory **1013** (e.g., a cache, registers, etc.). The programmable circuitry **1012** of the illustrated example is in communication with main memory **1014**, **1016**, which includes a volatile memory **1014** and a non-volatile memory **1016**, by a bus **1018**. The volatile memory **1014** may be implemented by Synchronous Dynamic Random Access Memory (SDRAM), Dynamic Random Access Memory (DRAM), RAMBUS® Dynamic Random Access Memory (RDRAM®), and/or any other type of RAM device. The non-volatile memory **1016** may be implemented by flash memory and/or any other desired type of memory device. Access to the main memory **1014**, **1016** of the illustrated example is controlled by a memory controller **1017**. In some examples, the memory controller **1017** may be implemented by one or more integrated circuits, logic circuits, microcontrollers from any desired family or manufacturer, or any other type of circuitry to manage the flow of data going to and from the main memory **1014**, **1016**.

[0093] The programmable circuitry platform **1000** of the illustrated example also includes interface circuitry **1020**. The interface circuitry **1020** may be implemented by hardware in accordance with any type of interface standard, such as an Ethernet interface, a universal serial bus (USB) interface, a Bluetooth® interface, a near field communication (NFC) interface, a Peripheral Component Interconnect (PCI) interface, and/or a Peripheral Component Interconnect Express (PCIe) interface. In some examples, the interface circuitry **1020** may include an output interface, such as an interface connected to a display device, an input interface such as an interface connected to an alphanumeric input device or a user interface (UI) navigation device, or a communication interface. In some examples, a connected I/O device may also include a display device, an alphanumeric input device, and/or a navigation device that is integrated into a single unit, such as a touch screen display. The communication interface may provide a connection with a network interface device used to transmit and/or receive electronic signals on the network **1026**. The programmable circuitry platform **1000** may also include other interfaces or hardware in connection with a signal generation device (e.g., an audio or radio signal generation device), an output controller (e.g., for connection with a serial, universal serial bus (USB), parallel, and/or other wired or wireless connection such as which uses via infrared (IR) and/or near field communication (NFC) technologies), an input controller (e.g., for connection with sensors or peripheral devices), etc. In some examples, the interface circuitry **1020** may implement the interface circuitry **300**, **400** of FIGS. **3** and/or **4**.

[0094] In the illustrated example, one or more input devices **1022** are connected to the interface circuitry **1020**. The input device(s) **1022** permit(s) a user (e.g., a human user, a machine user, etc.) to enter data and/or commands into the programmable circuitry **1012**. The input device(s) **1022** can be implemented by, for example, an audio sensor, a microphone, a camera (still or video), a keyboard, a button, a mouse, a touchscreen, and/or a voice recognition system.

[0095] One or more output devices **1024** are also connected to the interface circuitry **1020** of the illustrated example. The output device(s) **1024** can be implemented, for example, by display devices (e.g., a light emitting diode (LED), an organic light emitting diode (OLED), a liquid crystal display (LCD), a cathode ray tube (CRT) display, an in-place switching (IPS) display, a

touchscreen, etc.), and/or a tactile output device. The interface circuitry **1020** of the illustrated example, thus, typically includes a graphics driver card, a graphics driver chip, and/or graphics processor circuitry such as a GPU.

[0096] The interface circuitry **1020** of the illustrated example also includes a communication device such as a transmitter, a receiver, a transceiver, a modem, a residential gateway, a wireless access point, and/or a network interface to facilitate exchange of data with external machines (e.g., computing devices of any kind) by a network **1026**. The communication can be by, for example, an Ethernet connection, a digital subscriber line (DSL) connection, a telephone line connection, a coaxial cable system, a satellite system, a beyond-line-of-sight wireless system, a line-of-sight wireless system, a cellular telephone system, an optical connection, etc.

[0097] The programmable circuitry platform **1000** of the illustrated example also includes one or more mass storage discs or devices **1028** to store firmware, software, and/or data. Examples of such mass storage discs or devices **1028** include magnetic storage devices (e.g., floppy disk, drives, HDDs, etc.), optical storage devices (e.g., Blu-ray disks, CDs, DVDs, etc.), RAID systems, and/or solid-state storage discs or devices such as flash memory devices and/or SSDs.

[0098] The machine readable instructions **1032**, which may be implemented by the machine readable instructions of FIGS. **6** and/or **7**, may be stored in the mass storage device **1028**, in the volatile memory **1014**, in the non-volatile memory **1016**, and/or on at least one non-transitory computer readable storage medium such as a CD or DVD which may be removable. Some examples of a machine-readable medium are a non-transitory medium that hosts or stores one or more sets of data structures or instructions (e.g., software instructions) embodying or utilized by any one or more of the techniques or functions described herein. Such instructions are collectively labeled as instructions **1032**.

[0099] The instructions **1032** may reside, during execution and/or other operation of the programmable circuitry platform **1000**, completely, or at least partially, within the volatile memory **1014**, within non-volatile memory **1016**, within the local memory **1013**, within a removable storage, within a non-removable storage, and/or within the programmable circuitry **1012**. Thus, any combination of the programmable circuitry **1012**, the volatile memory **1014**, the non-volatile memory **1016**, the local memory **1013**, and/or a storage device of the removable storage or non-removable storage may constitute a machine-readable medium or media. The instructions **1032**, when loaded and executed by the programmable circuitry **1012**, may invoke or utilize a defined instruction set **1032** of the programmable circuitry **1012**, such as a processor instruction set defined by an instruction set architecture (ISA) of a reduced instruction set computer (RISC) or complex instruction set computer (CISC) architecture-including but not limited to the RISC-V Instruction Set provided in a RISC-V architecture. A RISC-V architecture and instruction set is one of several available architectures and instruction sets that may be used in examples of the compute components (e.g., the programmable circuitry **1012**) described herein.

[0100] FIG. **11** is a block diagram of an example implementation of the programmable circuitry **1012** of FIG. **10**. In this example, the programmable circuitry **1012** of FIG. **10** is implemented by a microprocessor **1100**. For example, the microprocessor **1100** may be a general-purpose microprocessor (e.g., general-purpose microprocessor circuitry). The microprocessor **1100** executes some or all of the machine-readable instructions of the flowcharts of FIGS. **6** and/or **7** to effectively instantiate the circuitry of FIGS. **2-4** as logic circuits to perform operations corresponding to those machine readable instructions. In some such examples, the circuitry of FIGS. **2-4** is instantiated by the hardware circuits of the microprocessor **1100** in combination with the machine-readable instructions. For example, the microprocessor **1100** may be implemented by multi-core hardware circuitry such as a CPU, a DSP, a GPU, a VPU, an XPU, etc. Although it may include any number of example cores **1102** (e.g., **1** core), the microprocessor **1100** of this example is a multi-core semiconductor device including N cores. The cores **1102** of the microprocessor **1100** may operate independently or may cooperate to execute machine readable instructions. For example, machine

code corresponding to a firmware program, an embedded software program, or a software program may be executed by one of the cores **1102** or may be executed by multiple ones of the cores **1102** at the same or different times. In some examples, the machine code corresponding to the firmware program, the embedded software program, or the software program is split into threads and executed in parallel by two or more of the cores **1102**. The software program may correspond to a portion or all of the machine readable instructions and/or operations represented by the flowcharts of FIGS. **6** and/or **7**.

[0101] The cores **1102** may communicate by a first example bus **1104**. In some examples, the first bus **1104** may be implemented by a communication bus to effectuate communication associated with one(s) of the cores **1102**. For example, the first bus **1104** may be implemented by at least one of an Inter-Integrated Circuit (I2C) bus, a Serial Peripheral Interface (SPI) bus, a PCI bus, or a PCIe bus. Additionally or alternatively, the first bus **1104** may be implemented by any other type of computing or electrical bus. The cores **1102** may obtain data, instructions, and/or signals from one or more external devices by example interface circuitry **1106**. The cores **1102** may output data, instructions, and/or signals to the one or more external devices by the interface circuitry **1106**. Although the cores **1102** of this example include example local memory **1120** (e.g., Level 1 (L1) cache that may be split into an L1 data cache and an L1 instruction cache), the microprocessor **1100** also includes example shared memory **1110** that may be shared by the cores (e.g., Level 2 (L2 cache)) for high-speed access to data and/or instructions. Data and/or instructions may be transferred (e.g., shared) by writing to and/or reading from the shared memory **1110**. The local memory **1120** of each of the cores **1102** and the shared memory **1110** may be part of a hierarchy of storage devices including multiple levels of cache memory and the main memory (e.g., the main memory **1014**, **1016** of FIG. **10**). Typically, higher levels of memory in the hierarchy exhibit lower access time and have smaller storage capacity than lower levels of memory. Changes in the various levels of the cache hierarchy are managed (e.g., coordinated) by a cache coherency policy.

[0102] Each core **1102** may be referred to as a CPU, DSP, GPU, etc., or any other type of hardware circuitry. Each core **1102** includes control unit circuitry **1114**, arithmetic and logic (AL) circuitry (sometimes referred to as an ALU) **1116**, a plurality of registers **1118**, the local memory **1120**, and a second example bus **1122**. Other structures may be present. For example, each core **1102** may include vector unit circuitry, single instruction multiple data (SIMD) unit circuitry, load/store unit (LSU) circuitry, branch/jump unit circuitry, floating-point unit (FPU) circuitry, etc. The control unit circuitry **1114** includes semiconductor-based circuits structured to control (e.g., coordinate) data movement within the corresponding core **1102**. The AL circuitry **1116** includes semiconductor-based circuits structured to perform one or more mathematic and/or logic operations on the data within the corresponding core **1102**. The AL circuitry **1116** of some examples performs integer based operations. In other examples, the AL circuitry **1116** also performs floating-point operations. In yet other examples, the AL circuitry **1116** may include first AL circuitry that performs integer-based operations and second AL circuitry that performs floating-point operations. In some examples, the AL circuitry **1116** may be referred to as an Arithmetic Logic Unit (ALU).

[0103] The registers **1118** are semiconductor-based structures to store data and/or instructions such as results of one or more of the operations performed by the AL circuitry **1116** of the corresponding core **1102**. For example, the registers **1118** may include vector register(s), SIMD register(s), general-purpose register(s), flag register(s), segment register(s), machine-specific register(s), instruction pointer register(s), control register(s), debug register(s), memory management register(s), machine check register(s), etc. The registers **1118** may be arranged in a bank as shown in FIG. **11**. Alternatively, the registers **1118** may be organized in any other arrangement, format, or structure, such as by being distributed throughout the core **1102** to shorten access time. The second bus **1122** may be implemented by at least one of an I2C bus, a SPI bus, a PCI bus, or a PCIe bus.

[0104] Each core **1102** and/or, more generally, the microprocessor **1100** may include additional and/or alternate structures to those shown and described above. For example, one or more clock

circuits, one or more power supplies, one or more power gates, one or more cache home agents (CHAs), one or more converged/common mesh stops (CMSs), one or more shifters (e.g., barrel shifter(s)) and/or other circuitry may be present. The microprocessor **1100** is a semiconductor device fabricated to include many transistors interconnected to implement the structures described above in one or more integrated circuits (ICs) contained in one or more packages.

[0105] The microprocessor **1100** may include and/or cooperate with one or more accelerators (e.g., acceleration circuitry, hardware accelerators, etc.). In some examples, accelerators are implemented by logic circuitry to perform certain tasks more quickly and/or efficiently than can be done by a general-purpose processor. Examples of accelerators include ASICs and FPGAs such as those discussed herein. A GPU, DSP and/or other programmable device can also be an accelerator. Accelerators may be on-board the microprocessor **1100**, in the same chip package as the microprocessor **1100** and/or in one or more separate packages from the microprocessor **1100**.

[0106] FIG. **12** is a block diagram of another example implementation of the programmable circuitry **1012** of FIG. **10**. In this example, the programmable circuitry **1012** is implemented by FPGA circuitry **1200**. For example, the FPGA circuitry **1200** may be implemented by an FPGA. The FPGA circuitry **1200** can be used, for example, to perform operations that could otherwise be performed by the example microprocessor **1100** of FIG. **11** executing corresponding machine readable instructions. However, once configured, the FPGA circuitry **1200** instantiates the operations and/or functions corresponding to the machine readable instructions in hardware and, thus, can often execute the operations/functions faster than they could be performed by a general-purpose microprocessor executing the corresponding software.

[0107] More specifically, in contrast to the microprocessor **1100** of FIG. **11** described above (which is a general purpose device that may be programmed to execute some or all of the machine readable instructions represented by the flowchart(s) of FIGS. **6** and/or **7** but whose interconnections and logic circuitry are fixed once fabricated), the FPGA circuitry **1200** of the example of FIG. **12** includes interconnections and logic circuitry that may be configured, structured, programmed, and/or interconnected in different ways after fabrication to instantiate, for example, some or all of the operations/functions corresponding to the machine readable instructions represented by the flowchart(s) of FIGS. **6** and/or **7**. In particular, the FPGA circuitry **1200** may be thought of as an array of logic gates, interconnections, and switches. The switches can be programmed to change how the logic gates are interconnected by the interconnections, effectively forming one or more dedicated logic circuits (unless and until the FPGA circuitry **1200** is reprogrammed). The configured logic circuits enable the logic gates to cooperate in different ways to perform different operations on data received by input circuitry. Those operations may correspond to some or all of the instructions (e.g., the software and/or firmware) represented by the flowchart(s) of FIGS. **6** and/or **7**. As such, the FPGA circuitry **1200** may be configured and/or structured to effectively instantiate some or all of the operations/functions corresponding to the machine readable instructions of the flowchart(s) of FIGS. **6** and/or **7** as dedicated logic circuits to perform the operations/functions corresponding to those software instructions in a dedicated manner analogous to an ASIC. Therefore, the FPGA circuitry **1200** may perform the operations/functions corresponding to the some or all of the machine readable instructions of FIGS. **6** and/or **7** faster than the general-purpose microprocessor can execute the same.

[0108] In the example of FIG. **12**, the FPGA circuitry **1200** is configured and/or structured in response to being programmed (and/or reprogrammed one or more times) based on a binary file. In some examples, the binary file may be compiled and/or generated based on instructions in a hardware description language (HDL) such as Lucid, Very High Speed Integrated Circuits (VHSIC) Hardware Description Language (VHDL), or Verilog. For example, a user (e.g., a human user, a machine user, etc.) may write code or a program corresponding to one or more operations/functions in an HDL; the code/program may be translated into a low-level language as needed; and the code/program (e.g., the code/program in the low-level language) may be converted

(e.g., by a compiler, a software application, etc.) into the binary file. In some examples, the FPGA circuitry **1200** of FIG. **12** may access and/or load the binary file to cause the FPGA circuitry **1200** of FIG. **12** to be configured and/or structured to perform the one or more operations/functions. For example, the binary file may be implemented by a bit stream (e.g., one or more computer-readable bits, one or more machine-readable bits, etc.), data (e.g., computer-readable data, machine-readable data, etc.), and/or machine-readable instructions accessible to the FPGA circuitry **1200** of FIG. **12** to cause configuration and/or structuring of the FPGA circuitry **1200** of FIG. **12**, or portion(s) thereof.

[0109] In some examples, the binary file is compiled, generated, transformed, and/or otherwise output from a uniform software platform utilized to program FPGAs. For example, the uniform software platform may translate first instructions (e.g., code or a program) that correspond to one or more operations/functions in a high-level language (e.g., C, C++, Python, etc.) into second instructions that correspond to the one or more operations/functions in an HDL. In some such examples, the binary file is compiled, generated, and/or otherwise output from the uniform software platform based on the second instructions. In some examples, the FPGA circuitry **1200** of FIG. **12** may access and/or load the binary file to cause the FPGA circuitry **1200** of FIG. **12** to be configured and/or structured to perform the one or more operations/functions. For example, the binary file may be implemented by a bit stream (e.g., one or more computer-readable bits, one or more machine-readable bits, etc.), data (e.g., computer-readable data, machine-readable data, etc.), and/or machine-readable instructions accessible to the FPGA circuitry **1200** of FIG. **12** to cause configuration and/or structuring of the FPGA circuitry **1200** of FIG. **12**, or portion(s) thereof.

[0110] The FPGA circuitry **1200** of FIG. **12**, includes example input/output (I/O) circuitry **1202** to obtain and/or output data to/from example configuration circuitry **1204** and/or external hardware **1206**. For example, the configuration circuitry **1204** may be implemented by interface circuitry that may obtain a binary file, which may be implemented by a bit stream, data, and/or machine-readable instructions, to configure the FPGA circuitry **1200**, or portion(s) thereof. In some such examples, the configuration circuitry **1204** may obtain the binary file from a user, a machine (e.g., hardware circuitry (e.g., programmable or dedicated circuitry) that may implement an Artificial Intelligence/Machine Learning (AI/ML) model to generate the binary file), etc., and/or any combination(s) thereof). In some examples, the external hardware **1206** may be implemented by external hardware circuitry. For example, the external hardware **1206** may be implemented by the microprocessor **1100** of FIG. **11**.

[0111] The FPGA circuitry **1200** also includes an array of example logic gate circuitry **1208**, a plurality of example configurable interconnections **1210**, and example storage circuitry **1212**. The logic gate circuitry **1208** and the configurable interconnections **1210** are configurable to instantiate one or more operations/functions that may correspond to at least some of the machine readable instructions of FIGS. **6** and/or **7** and/or other desired operations. The logic gate circuitry **1208** shown in FIG. **12** is fabricated in blocks or groups. Each block includes semiconductor-based electrical structures that may be configured into logic circuits. In some examples, the electrical structures include logic gates (e.g., And gates, Or gates, Nor gates, etc.) that provide basic building blocks for logic circuits. Electrically controllable switches (e.g., transistors) are present within each of the logic gate circuitry **1208** to enable configuration of the electrical structures and/or the logic gates to form circuits to perform desired operations/functions. The logic gate circuitry **1208** may include other electrical structures such as look-up tables (LUTs), registers (e.g., flip-flops or latches), multiplexers, etc.

[0112] The configurable interconnections **1210** of the illustrated example are conductive pathways, traces, vias, or the like that may include electrically controllable switches (e.g., transistors) whose state can be changed by programming (e.g., using an HDL instruction language) to activate or deactivate one or more connections between one or more of the logic gate circuitry **1208** to program desired logic circuits.

[0113] The storage circuitry **1212** of the illustrated example is structured to store result(s) of the one or more of the operations performed by corresponding logic gates. The storage circuitry **1212** may be implemented by registers or the like. In the illustrated example, the storage circuitry **1212** is distributed amongst the logic gate circuitry **1208** to facilitate access and increase execution speed.

[0114] The example FPGA circuitry **1200** of FIG. **12** also includes example dedicated operations circuitry **1214**. In this example, the dedicated operations circuitry **1214** includes special purpose circuitry **1216** that may be invoked to implement commonly used functions to avoid the need to program those functions in the field. Examples of such special purpose circuitry **1216** include memory (e.g., DRAM) controller circuitry, PCIe controller circuitry, clock circuitry, transceiver circuitry, memory, and multiplier-accumulator circuitry. Other types of special purpose circuitry may be present. In some examples, the FPGA circuitry **1200** may also include example general purpose programmable circuitry **1218** such as an example CPU **1220** and/or an example DSP **1222**. Other general purpose programmable circuitry **1218** may additionally or alternatively be present such as a GPU, an XPU, etc., that can be programmed to perform other operations.

[0115] Although FIGS. **11** and **12** illustrate two example implementations of the programmable circuitry **1012** of FIG. **10**, many other approaches are contemplated. For example, FPGA circuitry may include an on-board CPU, such as one or more of the example CPU **1220** of FIG. **11**. Therefore, the programmable circuitry **1012** of FIG. **10** may additionally be implemented by combining at least the example microprocessor **1100** of FIG. **11** and the example FPGA circuitry **1200** of FIG. **12**. In some such hybrid examples, one or more cores **1102** of FIG. **11** may execute a first portion of the machine readable instructions represented by the flowchart(s) of FIGS. **6** and/or **7** to perform first operation(s)/function(s), the FPGA circuitry **1200** of FIG. **12** may be configured and/or structured to perform second operation(s)/function(s) corresponding to a second portion of the machine readable instructions represented by the flowcharts of FIGS. **6** and/or **7**, and/or an ASIC may be configured and/or structured to perform third operation(s)/function(s) corresponding to a third portion of the machine readable instructions represented by the flowcharts of FIGS. **6** and/or **7**.

[0116] It should be understood that some or all of the circuitry of FIGS. **2-4** may, thus, be instantiated at the same or different times. For example, same and/or different portion(s) of the microprocessor **1100** of FIG. **11** may be programmed to execute portion(s) of machine-readable instructions at the same and/or different times. In some examples, same and/or different portion(s) of the FPGA circuitry **1200** of FIG. **12** may be configured and/or structured to perform operations/functions corresponding to portion(s) of machine-readable instructions at the same and/or different times.

[0117] In some examples, some or all of the circuitry of FIGS. **2-4** may be instantiated, for example, in one or more threads executing concurrently and/or in series. For example, the microprocessor **1100** of FIG. **11** may execute machine readable instructions in one or more threads executing concurrently and/or in series. In some examples, the FPGA circuitry **1200** of FIG. **12** may be configured and/or structured to carry out operations/functions concurrently and/or in series. Moreover, in some examples, some or all of the circuitry of FIGS. **2-4** may be implemented within one or more virtual machines and/or containers executing on the microprocessor **1100** of FIG. **11**.

[0118] In some examples, the programmable circuitry **1012** of FIG. **10** may be in one or more packages. For example, the microprocessor **1100** of FIG. **11** and/or the FPGA circuitry **1200** of FIG. **12** may be in one or more packages. In some examples, an XPU may be implemented by the programmable circuitry **1012** of FIG. **10**, which may be in one or more packages. For example, the XPU may include a CPU (e.g., the microprocessor **1100** of FIG. **11**, the CPU **1220** of FIG. **12**, etc.) in one package, a DSP (e.g., the DSP **1222** of FIG. **12**) in another package, a GPU in yet another package, and an FPGA (e.g., the FPGA circuitry **1200** of FIG. **12**) in still yet another package.

[0119] A block diagram illustrating an example software distribution platform **1305** to distribute software such as the example machine readable instructions **1032** of FIG. **10** to other hardware

devices (e.g., hardware devices owned and/or operated by third parties from the owner and/or operator of the software distribution platform) is illustrated in FIG. **13**. The example software distribution platform **1305** may be implemented by any computer server, data facility, cloud service, etc., capable of storing and transmitting software to other computing devices. The third parties may be customers of the entity owning and/or operating the software distribution platform **1305**. For example, the entity that owns and/or operates the software distribution platform **1305** may be a developer, a seller, and/or a licensor of software such as the example machine readable instructions **1032** of FIG. **10**. The third parties may be consumers, users, retailers, OEMs, etc., who purchase and/or license the software for use and/or re-sale and/or sub-licensing. In the illustrated example, the software distribution platform **1305** includes one or more servers and one or more storage devices. The storage devices store the machine readable instructions **1032**, which may correspond to the example machine readable instructions of FIGS. **6** and/or **7**, as described above. The one or more servers of the example software distribution platform **1305** are in communication with an example network **1310**, which may correspond to any one or more of the Internet and/or any of the example networks described above. In some examples, the one or more servers are responsive to requests to transmit the software to a requesting party as part of a commercial transaction. Payment for the delivery, sale, and/or license of the software may be handled by the one or more servers of the software distribution platform and/or by a third party payment entity. The servers enable purchasers and/or licensors to download the machine readable instructions **1032** from the software distribution platform **1305**. For example, the software, which may correspond to the example machine readable instructions of FIGS. **6** and/or **7**, may be downloaded to the example programmable circuitry platform **1000**, which is to execute the machine readable instructions **1032** to implement the example the thread(s) **106**, **108**, the compute element **110**, and/or the caching agent **114** of FIGS. **2**, **3**, and/or **4**. In some examples, one or more servers of the software distribution platform **1305** periodically offer, transmit, and/or force updates to the software (e.g., the example machine readable instructions **1032** of FIG. **10**) to ensure improvements, patches, updates, etc., are distributed and applied to the software at the end user devices. Although referred to as software above, the distributed "software" could alternatively be firmware.

[0120] The instructions **1032** may be transmitted or received over the network **1310** using a transmission medium via the interface circuitry **1020** of FIG. **10** and related devices utilizing any one of a number of transfer protocols (e.g., frame relay, internet protocol (IP), transmission control protocol (TCP), user datagram protocol (UDP), hypertext transfer protocol (HTTP), etc.). Example communication networks may include a local area network (LAN), a wide area network (WAN), a packet data network (e.g., the Internet), mobile telephone networks (e.g., cellular networks), and/or wireless data networks (e.g., Institute of Electrical and Electronics Engineers (IEEE) 802.11 family of standards known as Wi-Fi®), IEEE 802.15.4 family of standards, peer-to-peer (P2P) networks, among others.

[0121] A computing program may be written in any form of programming language, including compiled or interpreted languages, and it may be deployed in any form, including as a stand-alone program and/or as a module, component, subroutine, and/or other unit suitable for use in a computing environment. Also, programs, codes, and/or code segments for accomplishing the techniques described herein are construed as within the scope of the present disclosure by programmers of ordinary skill in the art.

[0122] "Including" and "comprising" (and all forms and tenses thereof) are used herein to be open ended terms. Thus, whenever a claim employs any form of "include" or "comprise" (e.g., comprises, includes, comprising, including, having, etc.) as a preamble or within a claim recitation of any kind, it is to be understood that additional elements, terms, etc., may be present without falling outside the scope of the corresponding claim or recitation. As used herein, when the phrase "at least" is used as the transition term in, for example, a preamble of a claim, it is open-ended in the same manner as the term "comprising" and "including" are open ended. The term "and/or"

when used, for example, in a form such as A, B, and/or C refers to any combination or subset of A, B, C such as (1) A alone, (2) B alone, (3) C alone, (4) A with B, (5) A with C, (6) B with C, or (7) A with B and with C. As used herein in the context of describing structures, components, items, objects and/or things, the phrase "at least one of A and B" is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B. Similarly, as used herein in the context of describing structures, components, items, objects and/or things, the phrase "at least one of A or B" is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B. As used herein in the context of describing the performance or execution of processes, instructions, actions, activities, etc., the phrase "at least one of A and B" is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B. Similarly, as used herein in the context of describing the performance or execution of processes, instructions, actions, activities, etc., the phrase "at least one of A or B" is intended to refer to implementations including any of (1) at least one A, (2) at least one B, or (3) at least one A and at least one B.

[0123] As used herein, singular references (e.g., "a," "an," "first," "second," etc.) do not exclude a plurality. The term "a" or "an" object, as used herein, refers to one or more of that object. The terms "a" (or "an"), "one or more," and "at least one" are used interchangeably herein. Furthermore, although individually listed, a plurality of means, elements, or actions may be implemented by, e.g., the same entity or object. Additionally, although individual features may be included in different examples or claims, these may possibly be combined, and the inclusion in different examples or claims does not imply that a combination of features is not feasible and/or advantageous.

[0124] As used herein, unless otherwise stated, the term "above" describes the relationship of two parts relative to Earth. A first part is above a second part, if the second part has at least one part between Earth and the first part. Likewise, as used herein, a first part is "below" a second part when the first part is closer to the Earth than the second part. As noted above, a first part can be above or below a second part with one or more of: other parts therebetween, without other parts therebetween, with the first and second parts touching, or without the first and second parts being in direct contact with one another.

[0125] As used in this patent, stating that any part (e.g., a layer, film, area, region, or plate) is in any way on (e.g., positioned on, located on, disposed on, or formed on, etc.) another part, indicates that the referenced part is either in contact with the other part, or that the referenced part is above the other part with one or more intermediate part(s) located therebetween.

[0126] As used herein, connection references (e.g., attached, coupled, connected, and joined) may include intermediate members between the elements referenced by the connection reference and/or relative movement between those elements unless otherwise indicated. As such, connection references do not necessarily infer that two elements are directly connected and/or in fixed relation to each other. As used herein, stating that any part is in "contact" with another part is defined to mean that there is no intermediate part between the two parts.

[0127] Unless specifically stated otherwise, descriptors such as "first," "second," "third," etc., are used herein without imputing or otherwise indicating any meaning of priority, physical order, arrangement in a list, and/or ordering in any way, but are merely used as labels and/or arbitrary names to distinguish elements for ease of understanding the disclosed examples. In some examples, the descriptor "first" may be used to refer to an element in the detailed description, while the same element may be referred to in a claim with a different descriptor such as "second" or "third." In such instances, it should be understood that such descriptors are used merely for identifying those elements distinctly within the context of the discussion (e.g., within a claim) in which the elements might, for example, otherwise share a same name.

[0128] As used herein, the phrase "in communication," including variations thereof, encompasses direct communication and/or indirect communication through one or more intermediary

components, and does not require direct physical (e.g., wired) communication and/or constant communication, but rather additionally includes selective communication at periodic intervals, scheduled intervals, aperiodic intervals, and/or one-time events.

[0129] As used herein, "programmable circuitry" is defined to include (i) one or more special purpose electrical circuits (e.g., an application specific circuit (ASIC)) structured to perform specific operation(s) and including one or more semiconductor-based logic devices (e.g., electrical hardware implemented by one or more transistors), and/or (ii) one or more general purpose semiconductor-based electrical circuits programmable with instructions to perform specific functions(s) and/or operation(s) and including one or more semiconductor-based logic devices (e.g., electrical hardware implemented by one or more transistors). Examples of programmable circuitry include programmable microprocessors such as Central Processor Units (CPUs) that may execute first instructions to perform one or more operations and/or functions, Field Programmable Gate Arrays (FPGAs) that may be programmed with second instructions to cause configuration and/or structuring of the FPGAs to instantiate one or more operations and/or functions corresponding to the first instructions, Graphics Processor Units (GPUs) that may execute first instructions to perform one or more operations and/or functions, Digital Signal Processors (DSPs) that may execute first instructions to perform one or more operations and/or functions, XPUs, Network Processing Units (NPUs) one or more microcontrollers that may execute first instructions to perform one or more operations and/or functions and/or integrated circuits such as Application Specific Integrated Circuits (ASICs). For example, an XPU may be implemented by a heterogeneous computing system including multiple types of programmable circuitry (e.g., one or more FPGAs, one or more CPUs, one or more GPUs, one or more NPUs, one or more DSPs, etc., and/or any combination(s) thereof), and orchestration technology (e.g., application programming interface(s) (API(s)) that may assign computing task(s) to whichever one(s) of the multiple types of programmable circuitry is/are suited and available to perform the computing task(s).

[0130] As used herein, integrated circuit/circuitry is defined as one or more semiconductor packages containing one or more circuit elements such as transistors, capacitors, inductors, resistors, current paths, diodes, etc. For example, an integrated circuit may be implemented as one or more of an ASIC, an FPGA, a chip, a microchip, programmable circuitry, a semiconductor substrate coupling multiple circuit elements, a system-on-chip (SoC), etc.

[0131] From the foregoing, it will be appreciated that example systems, apparatus, articles of manufacture, and methods have been disclosed that facilitate the use of dynamic coherency models for memory objects (e.g., memory objects that can be tagged for a coherency model or for consistency (e.g., no coherency)). Disclosed systems, apparatus, articles of manufacture, and methods improve the efficiency of using a computing device by dynamically subjecting memory objects to coherence when needed and removing coherence for the memory objects when not needed. In this manner, examples disclosed herein can reduce overhead and/or resources when coherency is not needed. Disclosed systems, apparatus, articles of manufacture, and methods are accordingly directed to one or more improvement(s) in the operation of a machine such as a computer or other electronic and/or mechanical device.

[0132] Example methods, apparatus, systems, and articles of manufacture to facilitate use of dynamic coherency models for memory objects are disclosed herein. Further examples and combinations thereof include the following:

[0133] Example 1 includes an apparatus comprising interface circuitry, instructions, and at least one programmable circuit to be programmed by the instructions to access a request identifying a memory range and a coherency model, and apply the coherency model to a memory address in the memory range based on the coherency model setting.

[0134] Example 2 includes the apparatus of one or more of the proceeding Examples, wherein one or more or the at least one programmable circuit is to configure a coherency setting for the memory range based on the request.

[0135] Example 3 includes the apparatus of one or more of the proceeding Examples, wherein the request includes a bitmask that identifies compute elements that are to implement the coherency model for the memory range.

[0136] Example 4 includes the apparatus of one or more of the proceeding Examples, wherein one or more of the at least one programmable circuit is to configure a coherency setting corresponding to the coherency model by storing an entry including the memory range and an identifier of the coherency model.

[0137] Example 5 includes the apparatus of one or more of the proceeding Examples, wherein one or more of the at least one programmable circuit is to configure a coherency setting corresponding to the coherency model by tagging the memory range.

[0138] Example 6 includes the apparatus of one or more of the proceeding Examples, wherein one or more of the at least one programmable circuit is to reset a coherency setting corresponding to the coherency model to a default coherency protocol, based on a reset coherency request.

[0139] Example 7 includes the apparatus of one or more of the proceeding Examples, wherein the interface circuitry is to receive the request from a compute element, the request generated based on a coherency instruction executed by the compute element.

[0140] Example 8 includes the apparatus of one or more of the proceeding Examples, wherein one or more of the at least one programmable circuit is to determine that a request to access the memory address corresponds to a coherency setting corresponding to the coherency model based on the memory address being within the memory range.

[0141] Example 9 includes a non-transitory computer readable medium comprising instructions which cause at least one programmable circuit to at least access a request identifying a memory range and a coherency model, and apply the coherency model to a memory address in the memory range based on the coherency model setting.

[0142] Example 10 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the instructions cause one or more of the at least one programmable circuit to configure a coherency setting for the memory range based on the request.

[0143] Example 11 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the request includes a bitmask that identifies processor elements that are to implement the coherency model for the memory range.

[0144] Example 12 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the instructions cause one or more of the at least one programmable circuit to configure a coherency setting corresponding to the coherency model by storing an entry including the memory range and an identifier of the coherency model.

[0145] Example includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the instructions cause one or more of the at least one programmable circuit to configure a coherency setting corresponding to the coherency model by tagging the memory range as coherent.

[0146] Example 14 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the instructions cause one or more of the at least one programmable circuit to reset a coherency setting corresponding to the coherency model to a default coherency protocol based on a revert coherency request.

[0147] Example 15 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the instructions cause one or more of the at least one programmable circuit to receive the request from a processor element, the request generated based on a coherency instruction executed by the processor element.

[0148] Example 16 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the instructions cause one or more of the at least one programmable circuit to determine that a request to access the memory address corresponds to a coherency setting corresponding to the coherency model based on the memory address being within the memory

range.

[0149] Example 17 includes a method comprising accessing a request identifying a memory range and a coherency model, and applying the coherency model to a memory address in the memory range based on the coherency model setting.

[0150] Example 18 includes the method of one or more of the proceeding Examples, including configuring a coherency setting for the memory range based on the request.

[0151] Example 19 includes the method of one or more of the proceeding Examples, wherein the request includes a bitmask that identifies central processing units that are to implement the coherency model for the memory range.

[0152] Example 20 includes the method of one or more of the proceeding Examples, including configuring a coherency setting corresponding to the coherency model includes storing an entry including the memory range and an identifier of the coherency model.

[0153] Example 21 includes the method of one or more of the proceeding Examples, including configuring a coherency setting corresponding to the coherency model includes tagging the memory range as coherent.

[0154] Example 22 includes the method of one or more of the proceeding Examples, including resetting a coherency setting corresponding to the coherency model to a default coherency protocol based on a reset coherency request.

[0155] Example 23 includes the method of one or more of the proceeding Examples, including receiving the request and from a central processing unit, the request generated based on a coherency instruction executed by the central processing unit.

[0156] Example 24 includes the method of one or more of the proceeding Examples, including determining that a request to access the memory address corresponds to a coherency setting corresponding to the coherency model based on the memory address being within the memory range.

[0157] Example 25 includes a system comprising a compute element to transmit a first request based on an instruction to implement a coherency model for a memory range, and transmit a second request to access a memory address based on an instruction to perform an operation with data corresponding to the memory address, and a caching agent to configure a coherency setting for the memory range based on the first request, determine if the memory address included in the second request is within the memory range, and facilitate the coherency model for the second request based on the memory address being included in the memory range.

[0158] Example 26 includes the system of one or more of the proceeding Examples, wherein the memory address is a first memory address, the caching agent to determine whether a second memory address included in a third request is within the memory range, and grant permission to the third request based on the second memory address being outside the memory range.

[0159] Example 27 includes a non-transitory computer readable medium comprising instructions which cause at least one programmable circuit to at least transmit a first request based on an instruction to implement coherency for a memory range, transmit a second request to access a memory address based on an instruction to perform an operation with data corresponding to the memory address, configure a coherency setting for the memory range based on the first request, determine if the memory address included in the second request is within the memory range, and facilitate a coherency model for the second request based on the memory address being included in the memory range.

[0160] Example 28 includes the non-transitory computer readable medium of one or more of the proceeding Examples, wherein the memory address is a first memory address, the instructions to cause one or more of the at least one programmable circuit to determine whether a second memory address included in a third request is within the memory range, and grant permission to the third request based on the second memory address being outside the memory range.

[0161] Example 29 includes a method comprising transmitting a first request based on an

instruction to implement coherency for a memory range, transmitting an second request for a memory address based on an instruction to perform an operation with data corresponding to the memory address, a coherency setting for the memory range based on the first request, determining if the memory address included in the second request is within the memory range, and facilitating a coherency model for the second request based on the memory address being included in the memory range.

[0162] Example 30 includes the method of one or more of the proceeding Examples, wherein the memory address is a first memory address, including determining whether a second memory address included in a third request is within the memory range, and granting permission to the third request based on the second memory address being outside the memory range.

[0163] Example 31 includes an apparatus comprising means for accessing a request identifying a memory range and a coherency model, and means for applying the coherency model to a memory address in the memory range based on the coherency model setting.

[0164] Example 32 includes the apparatus of one or more of the proceeding Examples, including means for configuring a coherency setting for the memory range based on the request.

[0165] Example 33 includes the apparatus of one or more of the proceeding Examples, wherein the request includes a bitmask that identifies compute elements that are to implement the coherency model for the memory range.

[0166] Example 34 includes the apparatus of one or more of the proceeding Examples, including means for configuring a coherency setting corresponding to the coherency model by storing an entry including the memory range and an identifier of the coherency model.

[0167] Example 35 includes the apparatus of one or more of the proceeding Examples, including means for configuring a coherency setting corresponding to the coherency model by tagging the memory range.

[0168] Example 36 includes the apparatus of one or more of the proceeding Examples, wherein the means for applying is to reset a coherency setting corresponding to the coherency model to a default coherency protocol, based on a reset coherency request.

[0169] Example 37 includes the apparatus of one or more of the proceeding Examples, wherein the means for accessing is to access the request from a compute element, the request generated based on a coherency instruction executed by the compute element.

[0170] Example 38 includes the apparatus of one or more of the proceeding Examples, including means for determining that a request to access the memory address corresponds to a coherency setting corresponding to the coherency model based on the memory address being within the memory range.

[0171] Example 39 includes a system comprising means for transmitting to transmit a first request based on an instruction to implement a coherency model for a memory range, and transmit a second request to access a memory address based on an instruction to perform an operation with data corresponding to the memory address, and means for configuring to configure a coherency setting for the memory range based on the first request, determine if the memory address included in the second request is within the memory range, and facilitate the coherency model for the second request based on the memory address being included in the memory range.

[0172] Example 40 includes the system of one or more of the proceeding Examples, wherein the memory address is a first memory address, the means for configuring to determine whether a second memory address included in a third request is within the memory range, and grant permission to the third request based on the second memory address being outside the memory range.

[0173] Example 41 includes machine-readable storage including machine-readable instructions, when executed, to implement a method as claimed in any of the examples 17 to 24 and 29 to 30.

[0174] The following claims are hereby incorporated into this Detailed Description by this reference. Although certain example systems, apparatus, articles of manufacture, and methods have

been disclosed herein, the scope of coverage of this patent is not limited thereto. On the contrary, this patent covers all systems, apparatus, articles of manufacture, and methods fairly falling within the scope of the claims of this patent.

Claims

1. An apparatus comprising: interface circuitry; instructions; and at least one programmable circuit to be programmed by the instructions to: access a request identifying a memory range and a coherency model; and apply the coherency model to a memory address in the memory range based on a coherency setting.

2. The apparatus of claim 1, wherein one or more or the at least one programmable circuit is to configure the coherency setting for the memory range based on the request.

3. The apparatus of claim 1, wherein the request includes a bitmask that identifies compute elements that are to implement the coherency model for the memory range.

4. The apparatus of claim 1, wherein one or more of the at least one programmable circuit is to configure the coherency setting corresponding to the coherency model by storing an entry including the memory range and an identifier of the coherency model.

5. The apparatus of claim 1, wherein one or more of the at least one programmable circuit is to configure the coherency setting corresponding to the coherency model by tagging the memory range.

6. The apparatus of claim 1, wherein one or more of the at least one programmable circuit is to reset the coherency setting corresponding to the coherency model to a default coherency protocol, based on a reset coherency request.

7. The apparatus of claim 1, wherein the interface circuitry is to receive the request from a compute element, the request generated based on a coherency instruction executed by the compute element.

8. The apparatus of claim 1, wherein one or more of the at least one programmable circuit is to determine that a request to access the memory address corresponds to the coherency setting corresponding to the coherency model based on the memory address being within the memory range.

9. A non-transitory computer readable medium comprising instructions which cause at least one programmable circuit to at least: access a request identifying a memory range and a coherency model; and apply the coherency model to a memory address in the memory range based on a coherency model setting.

10. The non-transitory computer readable medium of claim 9, wherein the instructions cause one or more of the at least one programmable circuit to configure the coherency setting for the memory range based on the request.

11. The non-transitory computer readable medium of claim 9, wherein the request includes a bitmask that identifies processor elements that are to implement the coherency model for the memory range.

12. The non-transitory computer readable medium of claim 9, wherein the instructions cause one or more of the at least one programmable circuit to configure the coherency setting corresponding to the coherency model by storing an entry including the memory range and an identifier of the coherency model.

13. The non-transitory computer readable medium of claim 9, wherein the instructions cause one or more of the at least one programmable circuit to configure the coherency setting corresponding to the coherency model by tagging the memory range as coherent.

14. The non-transitory computer readable medium of claim 9, wherein the instructions cause one or more of the at least one programmable circuit to reset the coherency setting corresponding to the coherency model to a default coherency protocol based on a revert coherency request.

15. The non-transitory computer readable medium of claim 9, wherein the instructions cause one or

more of the at least one programmable circuit to receive the request from a processor element, the request generated based on a coherency instruction executed by the processor element.

16. The non-transitory computer readable medium of claim 9, wherein the instructions cause one or more of the at least one programmable circuit to determine that a request to access the memory address corresponds to the coherency setting corresponding to the coherency model based on the memory address being within the memory range.

17.-**30**. (canceled)

31. An apparatus comprising: means for accessing a request identifying a memory range and a coherency model; and means for applying the coherency model to a memory address in the memory range based on a coherency setting.

32. The apparatus of claim 31, including means for configuring the coherency setting for the memory range based on the request.

33. The apparatus of claim 31, wherein the request includes a bitmask that identifies compute elements that are to implement the coherency model for the memory range.

34. The apparatus of one of claim 31, including means for configuring the coherency setting corresponding to the coherency model by storing an entry including the memory range and an identifier of the coherency model.

35.-**41**. (canceled)