



US 20250258796A1

(19) **United States**

(12) **Patent Application Publication**
Akidau et al.

(10) **Pub. No.: US 2025/0258796 A1**

(43) **Pub. Date: Aug. 14, 2025**

(54) **CONTINUOUS INGESTION OF CUSTOM
FILE FORMATS**

Publication Classification

(71) Applicant: **Snowflake Inc.**, Bozeman, MT (US)

(72) Inventors: **Tyler Arthur Akidau**, Seattle, WA
(US); **Thierry Cruanes**, San Mateo,
CA (US); **Benoit Dageville**, San Mateo,
CA (US); **Ganeshan Ramachandran**
Iyer, Redmond, WA (US);
Subramanian Muralidhar, Mercer
Island, WA (US)

(51) **Int. Cl.**

G06F 16/14 (2019.01)

G06F 9/50 (2006.01)

G06F 16/11 (2019.01)

(52) **U.S. Cl.**

CPC **G06F 16/148** (2019.01); **G06F 9/5022**
(2013.01); **G06F 16/116** (2019.01)

(21) Appl. No.: **19/193,450**

(22) Filed: **Apr. 29, 2025**

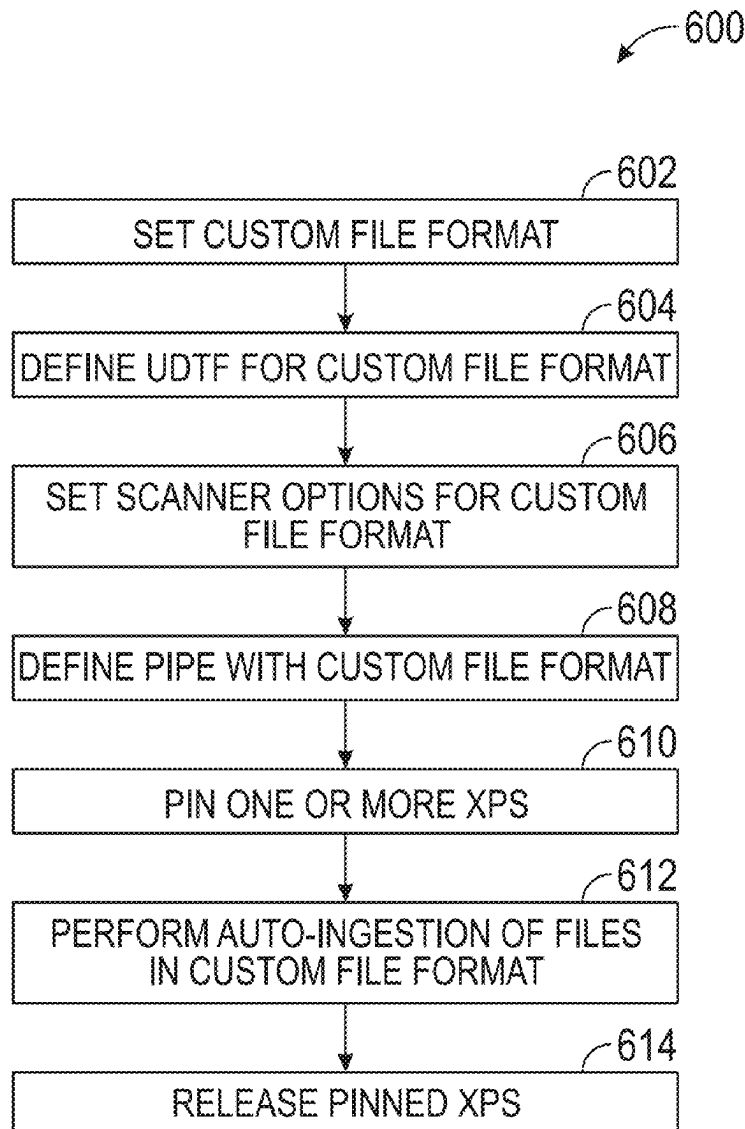
Related U.S. Application Data

(63) Continuation of application No. 18/050,122, filed on
Oct. 27, 2022, now Pat. No. 12,314,225.

(57)

ABSTRACT

Techniques for continuous ingestion of files using custom file formats are described. A custom file format may include formats not natively supported by a data system. Unstructured files (e.g., images) may also be considered custom file formats. A custom file format may be set using a user defined table function and scanner options.



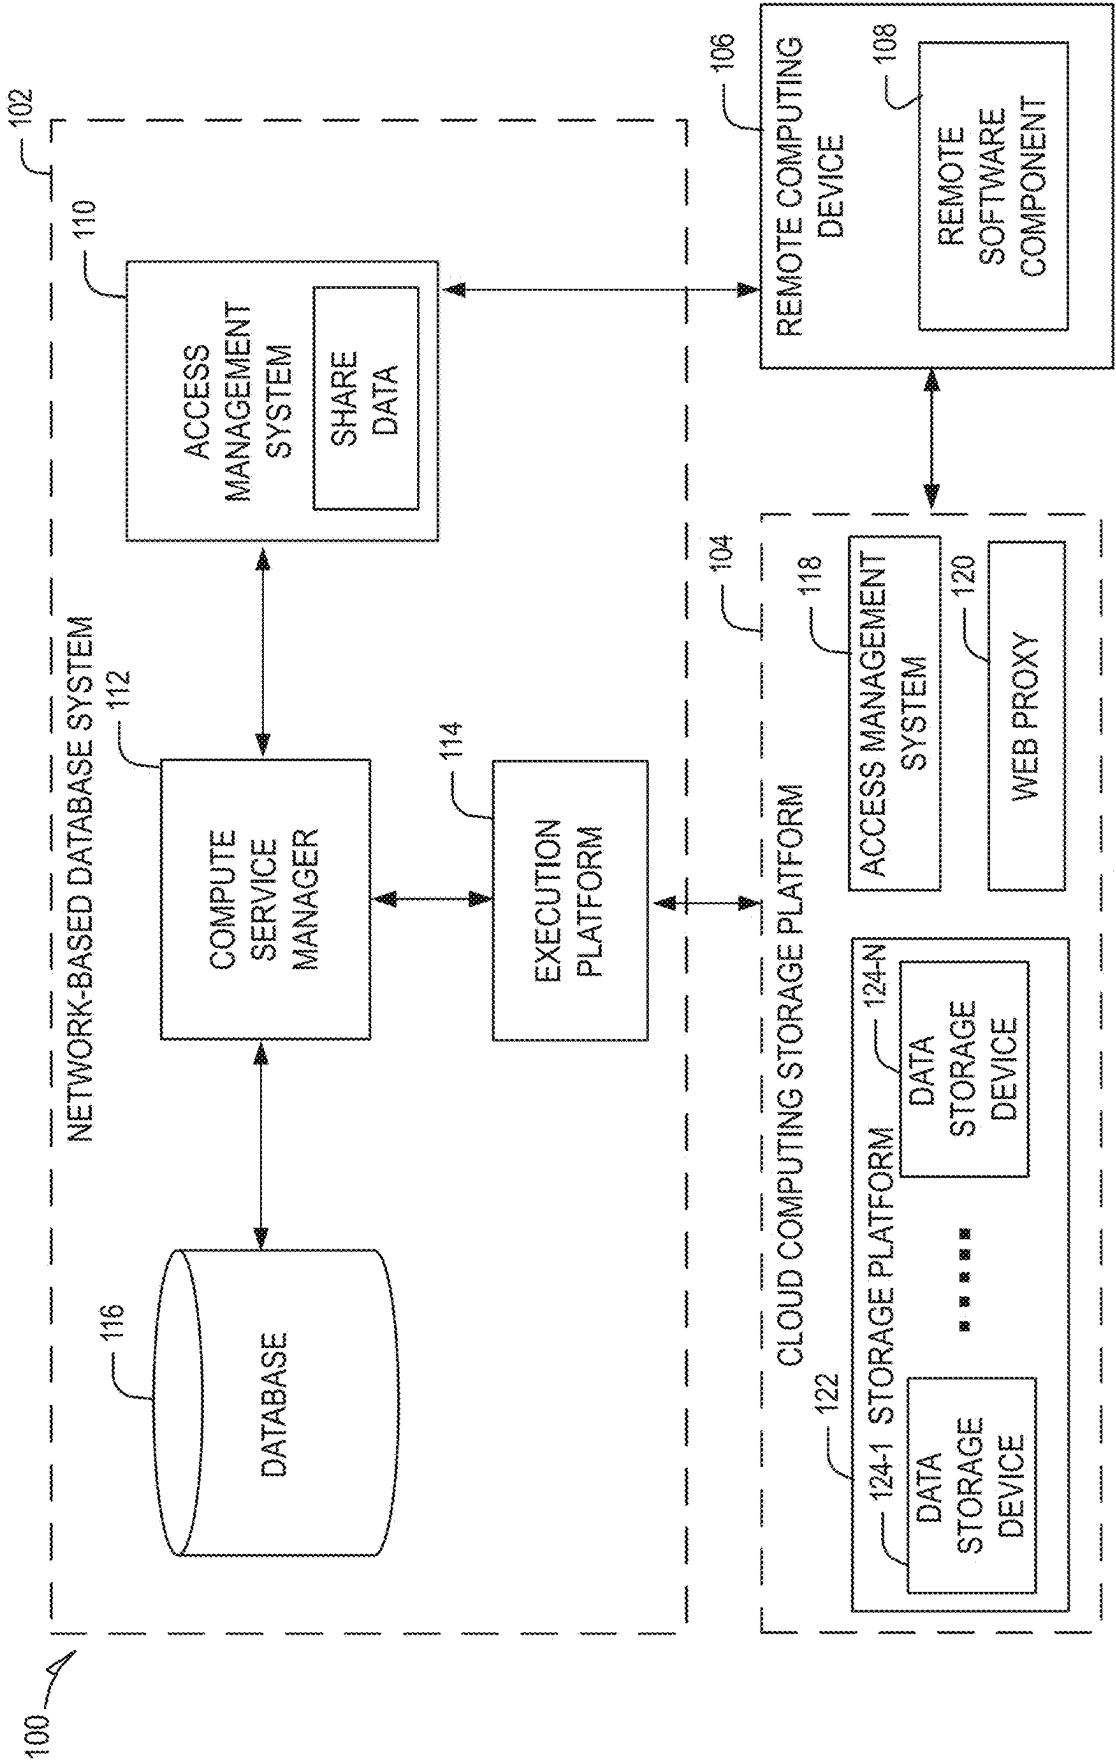


FIG. 1

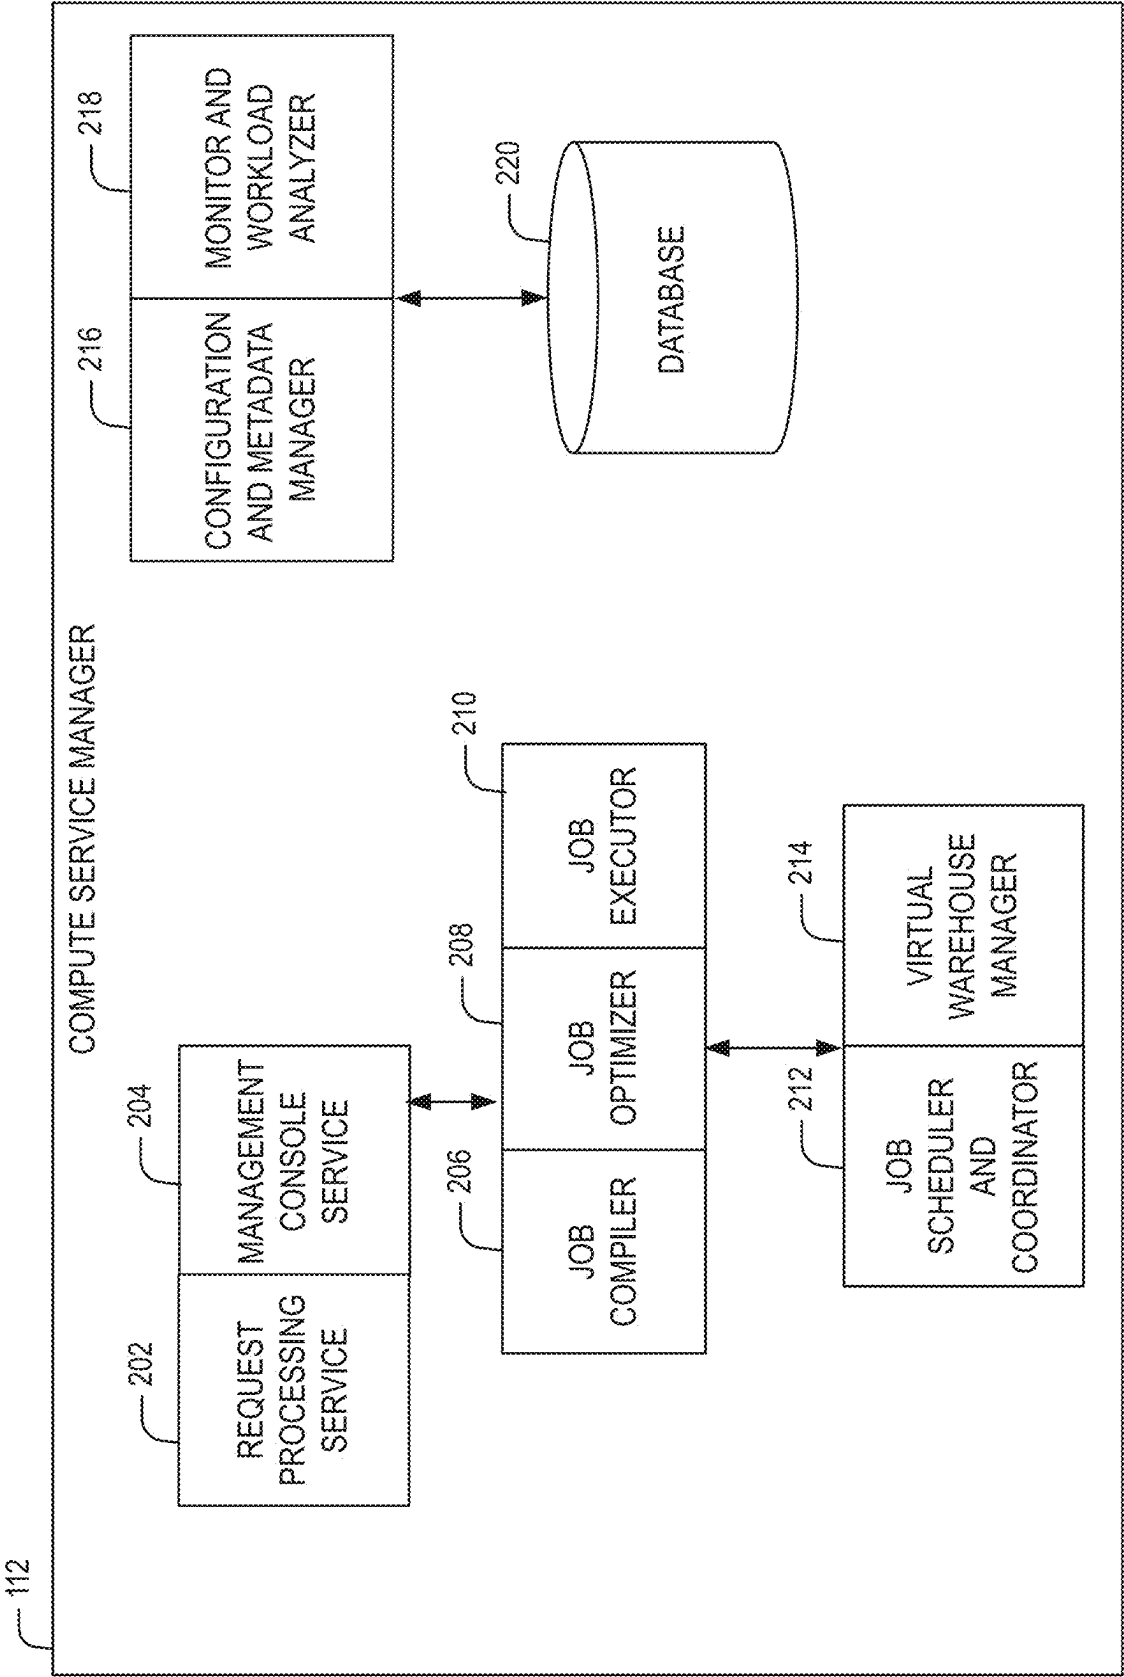


FIG. 2




FIG. 3



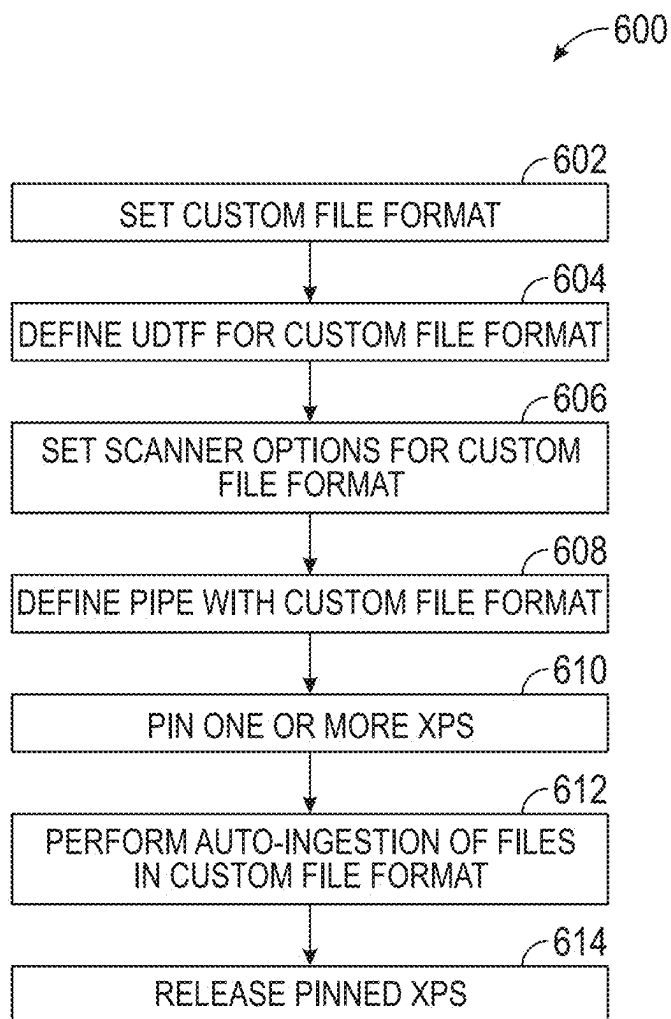


FIG. 4



**FIG. 6**

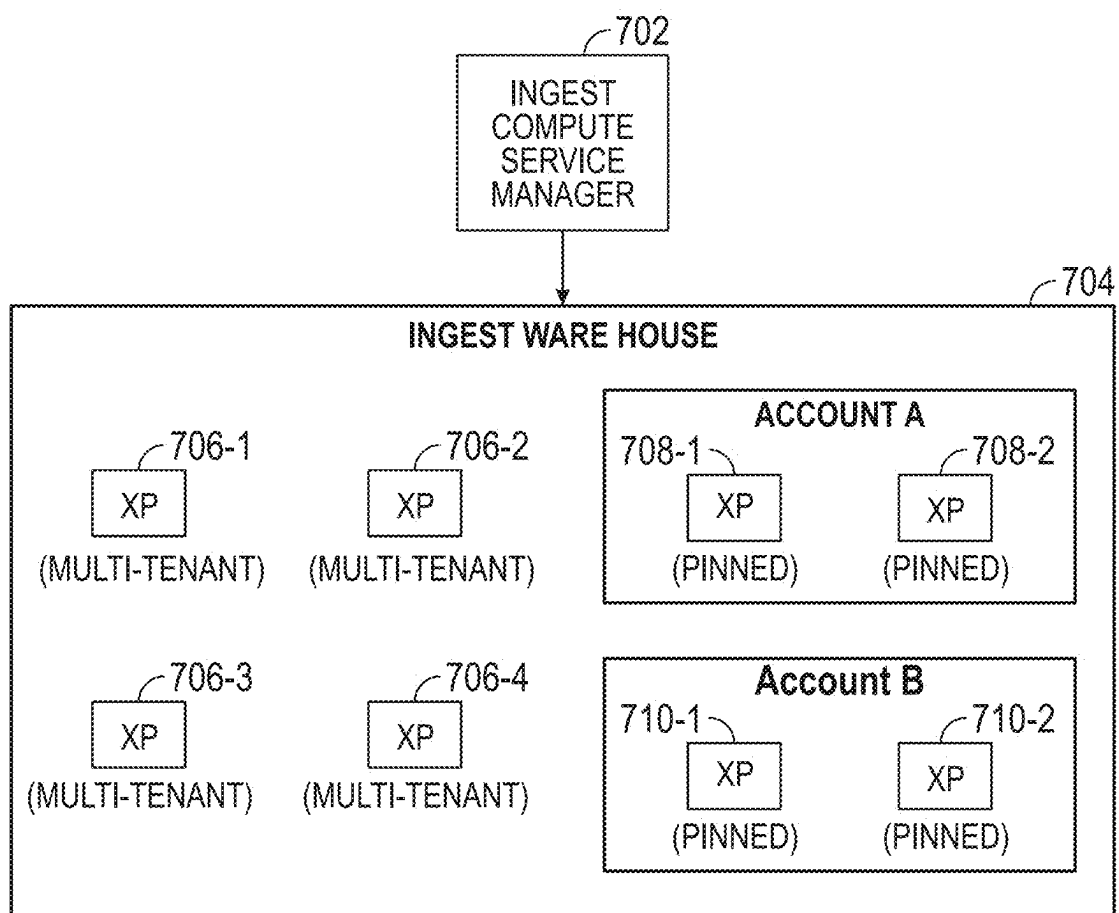


FIG. 7

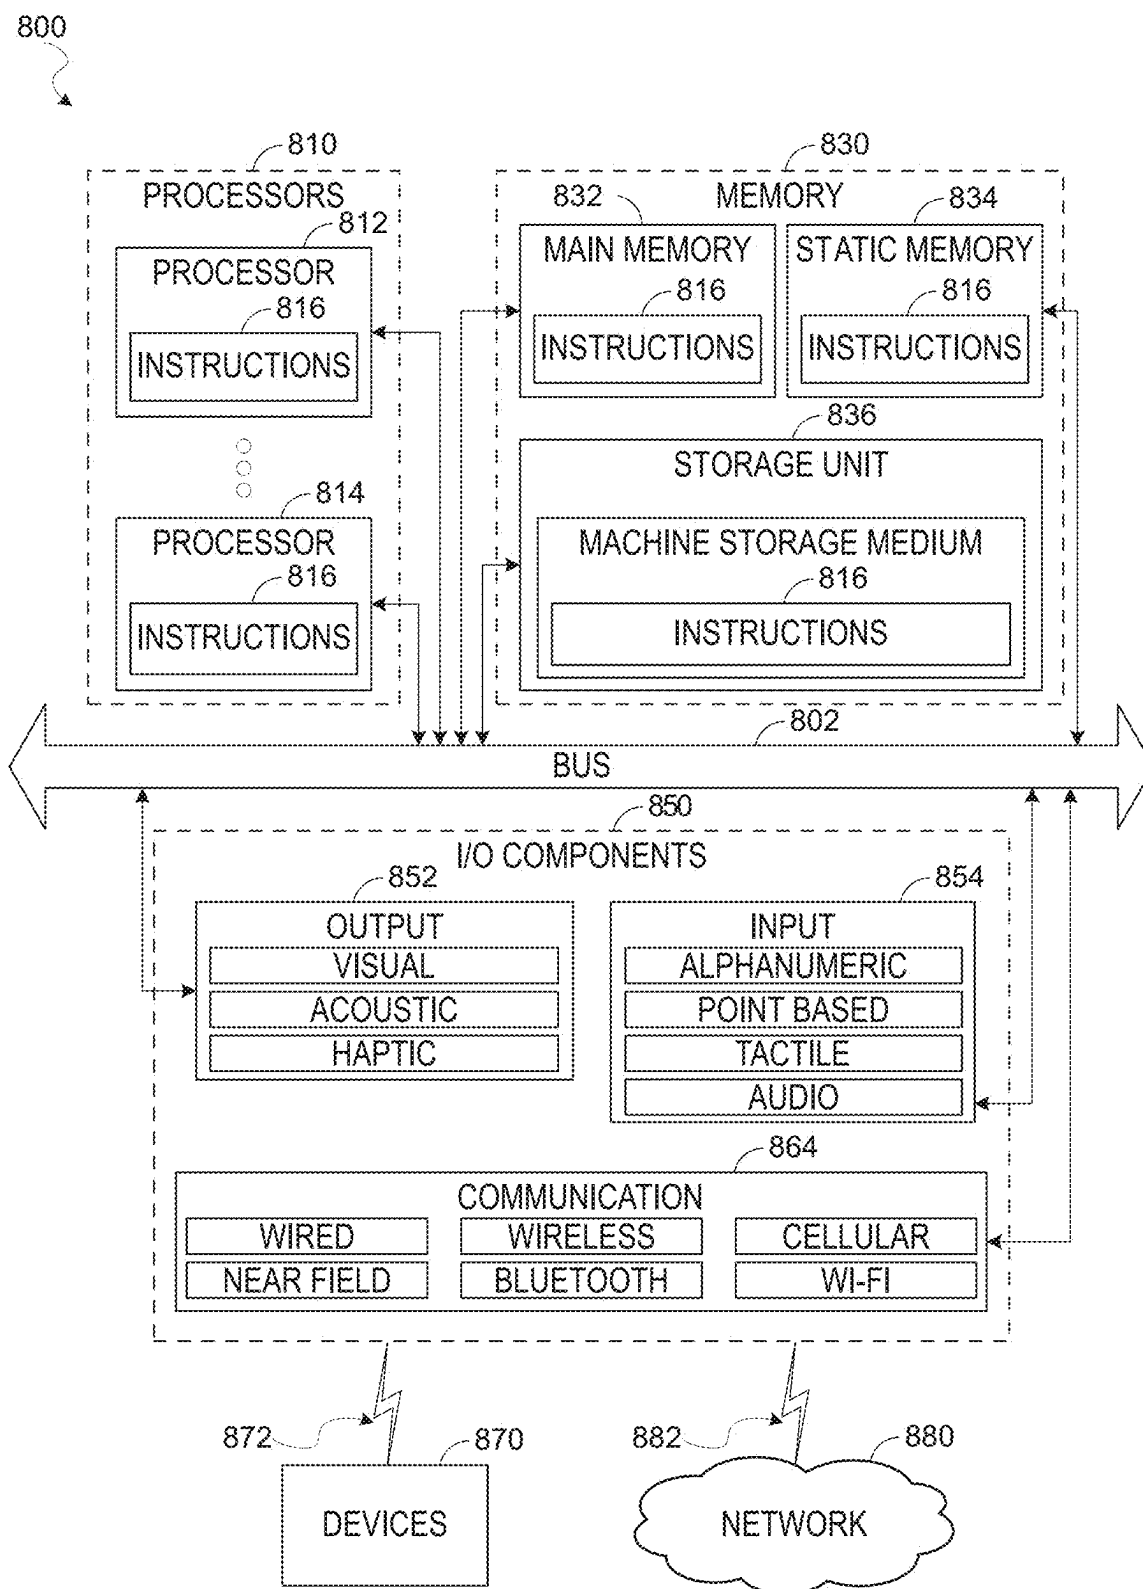


FIG. 8

CONTINUOUS INGESTION OF CUSTOM FILE FORMATS

PRIORITY CLAIM

[0001] This application is a Continuation of U.S. patent application Ser. No. 18/050,122, filed Oct. 27, 2022, the content of which is hereby incorporated by reference in its entirety.

TECHNICAL FIELD

[0002] The present disclosure generally relates to data systems, such as data systems, and, more specifically, to data ingestion techniques for different file formats.

BACKGROUND

[0003] Data systems, such as database systems, may be provided through a cloud platform, which allows organizations and users to store, manage, and retrieve data from the cloud. A variety of techniques can be employed for uploading and storing data in a database or table in a cloud platform. Uploading techniques typically cannot account for different file formats.

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] Various ones of the appended drawings merely illustrate example embodiments of the present disclosure and should not be considered as limiting its scope.

[0005] FIG. 1 illustrates an example computing environment with a cloud database system, according to some example embodiments.

[0006] FIG. 2 is a block diagram illustrating components of a compute service manager, according to some example embodiments.

[0007] FIG. 3 is a block diagram illustrating components of an execution platform, according to some example embodiments.

[0008] FIG. 4 is a simplified block diagram of a system for automated data ingestion, according to some example embodiments.

[0009] FIG. 5 is a schematic block diagram of a process of ingesting data into a database, according to some example embodiments.

[0010] FIG. 6 is a flow diagram of a method for auto-ingestion of custom file formats, according to some example embodiments.

[0011] FIG. 7 illustrates a data system framework for ingestion of files with custom file formats, according to some example embodiments.

[0012] FIG. 8 illustrates a diagrammatic representation of a machine in the form of a computer system within which a set of instructions may be executed for causing the machine to perform any one or more of the methodologies discussed herein, in accordance with some embodiments of the present disclosure.

DETAILED DESCRIPTION

[0013] The description that follows includes systems, methods, techniques, instruction sequences, and computing machine program products that embody illustrative embodiments of the disclosure. In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide an understanding of various

embodiments of the inventive subject matter. It will be evident, however, to those skilled in the art, that embodiments of the inventive subject matter may be practiced without these specific details. In general, well-known instruction instances, protocols, structures, and techniques are not necessarily shown in detail.

[0014] Techniques for continuous ingestion of files using custom file formats are described. A custom file format may include formats not natively supported by a data system. Unstructured files (e.g., images) may also be considered custom file formats.

[0015] FIG. 1 illustrates an example shared data processing platform 100. To avoid obscuring the inventive subject matter with unnecessary detail, various functional components that are not germane to conveying an understanding of the inventive subject matter have been omitted from the figures. However, a skilled artisan will readily recognize that various additional functional components may be included as part of the shared data processing platform 100 to facilitate additional functionality that is not specifically described herein.

[0016] As shown, the shared data processing platform 100 comprises the network-based database system 102, a cloud computing storage platform 104 (e.g., a storage platform, an AWS® service, Microsoft Azure®, or Google Cloud Services®), and a remote computing device 106. The network-based database system 102 is a cloud database system used for storing and accessing data (e.g., internally storing data, accessing external remotely located data) in an integrated manner, and reporting and analysis of the integrated data from the one or more disparate sources (e.g., the cloud computing storage platform 104). The cloud computing storage platform 104 comprises a plurality of computing machines and provides on-demand computer system resources such as data storage and computing power to the network-based database system 102. While in the embodiment illustrated in FIG. 1, a data warehouse is depicted, other embodiments may include other types of databases or other data processing systems.

[0017] The remote computing device 106 (e.g., a user device such as a laptop computer) comprises one or more computing machines (e.g., a user device such as a laptop computer) that execute a remote software component 108 (e.g., browser accessed cloud service) to provide additional functionality to users of the network-based database system 102. The remote software component 108 comprises a set of machine-readable instructions (e.g., code) that, when executed by the remote computing device 106, cause the remote computing device 106 to provide certain functionality. The remote software component 108 may operate on input data and generates result data based on processing, analyzing, or otherwise transforming the input data. As an example, the remote software component 108 can be a data provider or data consumer that enables database tracking procedures.

[0018] The network-based database system 102 comprises an access management system 110, a compute service manager 112, an execution platform 114, and a database 116. The access management system 110 enables administrative users to manage access to resources and services provided by the network-based database system 102. Administrative users can create and manage users, roles, and groups, and use permissions to allow or deny access to resources and services. The access management system 110 can store shared

data that securely manages shared access to the storage resources of the cloud computing storage platform **104** amongst different users of the network-based database system **102**, as discussed in further detail below.

[0019] The compute service manager **112** coordinates and manages operations of the network-based database system **102**. The compute service manager **112** also performs query optimization and compilation as well as managing clusters of computing services that provide compute resources (e.g., virtual warehouses, virtual machines, EC2 clusters). The compute service manager **112** can support any number of client accounts such as end users providing data storage and retrieval requests, system administrators managing the systems and methods described herein, and other components/devices that interact with compute service manager **112**.

[0020] The compute service manager **112** is also coupled to database **116**, which is associated with the entirety of data stored on the shared data processing platform **100**. The database **116** stores data pertaining to various functions and aspects associated with the network-based database system **102** and its users.

[0021] In some embodiments, database **116** includes a summary of data stored in remote data storage systems as well as data available from one or more local caches. Additionally, database **116** may include information regarding how data is organized in the remote data storage systems and the local caches. Database **116** allows systems and services to determine whether a piece of data needs to be accessed without loading or accessing the actual data from a storage device. The compute service manager **112** is further coupled to an execution platform **114**, which provides multiple computing resources (e.g., virtual warehouses) that execute various data storage and data retrieval, as discussed in greater detail below.

[0022] Execution platform **114** is coupled to multiple data storage devices **124-1** to **124-N** that are part of a cloud computing storage platform **104**. In some embodiments, data storage devices **124-1** to **124-N** are cloud-based storage devices located in one or more geographic locations. For example, data storage devices **124-1** to **124-N** may be part of a public cloud infrastructure or a private cloud infrastructure. Data storage devices **124-1** to **124-N** may be hard disk drives (HDDs), solid state drives (SSDs), storage clusters, Amazon S3 storage systems or any other data storage technology. Additionally, cloud computing storage platform **104** may include distributed file systems (such as Hadoop Distributed File Systems (HDFS)), object storage systems, and the like.

[0023] The execution platform **114** comprises a plurality of compute nodes (e.g., virtual warehouses). A set of processes on a compute node executes a query plan compiled by the compute service manager **112**. The set of processes can include: a first process to execute the query plan; a second process to monitor and delete micro-partition files using a least recently used (LRU) policy, and implement an out of memory (OOM) error mitigation process; a third process that extracts health information from process logs and status information to send back to the compute service manager **112**; a fourth process to establish communication with the compute service manager **112** after a system boot; and a fifth process to handle all communication with a compute cluster for a given job provided by the compute service manager

112 and to communicate information back to the compute service manager **112** and other compute nodes of the execution platform **114**.

[0024] The cloud computing storage platform **104** also comprises an access management system **118** and a web proxy **120**. As with the access management system **110**, the access management system **118** allows users to create and manage users, roles, and groups, and use permissions to allow or deny access to cloud services and resources. The access management system **110** of the network-based database system **102** and the access management system **118** of the cloud computing storage platform **104** can communicate and share information so as to enable access and management of resources and services shared by users of both the network-based database system **102** and the cloud computing storage platform **104**. The web proxy **120** handles tasks involved in accepting and processing concurrent API calls, including traffic management, authorization and access control, monitoring, and API version management. The web proxy **120** provides HTTP proxy service for creating, publishing, maintaining, securing, and monitoring APIs (e.g., REST APIs).

[0025] In some embodiments, communication links between elements of the shared data processing platform **100** are implemented via one or more data communication networks. These data communication networks may utilize any communication protocol and any type of communication medium. In some embodiments, the data communication networks are a combination of two or more data communication networks (or sub-Networks) coupled to one another. In alternative embodiments, these communication links are implemented using any type of communication medium and any communication protocol.

[0026] As shown in FIG. 1, data storage devices **124-1** to **124-N** are decoupled from the computing resources associated with the execution platform **114**. That is, new virtual warehouses can be created and terminated in the execution platform **114** and additional data storage devices can be created and terminated on the cloud computing storage platform **104** in an independent manner. This architecture supports dynamic changes to the network-based database system **102** based on the changing data storage/retrieval needs as well as the changing needs of the users and systems accessing the shared data processing platform **100**. The support of dynamic changes allows network-based database system **102** to scale quickly in response to changing demands on the systems and components within network-based database system **102**. The decoupling of the computing resources from the data storage devices **124-1** to **124-N** supports the storage of large amounts of data without requiring a corresponding large amount of computing resources. Similarly, this decoupling of resources supports a significant increase in the computing resources utilized at a particular time without requiring a corresponding increase in the available data storage resources. Additionally, the decoupling of resources enables different accounts to handle creating additional compute resources to process data shared by other users without affecting the other users' systems. For instance, a data provider may have three compute resources and share data with a data consumer, and the data consumer may generate new compute resources to execute queries against the shared data, where the new compute resources are managed by the data consumer and do not affect or interact with the compute resources of the data provider.

[0027] Compute service manager 112, database 116, execution platform 114, cloud computing storage platform 104, and remote computing device 106 are shown in FIG. 1 as individual components. However, each of compute service manager 112, database 116, execution platform 114, cloud computing storage platform 104, and remote computing environment may be implemented as a distributed system (e.g., distributed across multiple systems/platforms at multiple geographic locations) connected by APIs and access information (e.g., tokens, login data). Additionally, each of compute service manager 112, database 116, execution platform 114, and cloud computing storage platform 104 can be scaled up or down (independently of one another) depending on changes to the requests received and the changing needs of shared data processing platform 100. Thus, in the described embodiments, the network-based database system 102 is dynamic and supports regular changes to meet the current data processing needs.

[0028] During typical operation, the network-based database system 102 processes multiple jobs (e.g., queries) determined by the compute service manager 112. These jobs are scheduled and managed by the compute service manager 112 to determine when and how to execute the job. For example, the compute service manager 112 may divide the job into multiple discrete tasks and may determine what data is needed to execute each of the multiple discrete tasks. The compute service manager 112 may assign each of the multiple discrete tasks to one or more nodes of the execution platform 114 to process the task. The compute service manager 112 may determine what data is needed to process a task and further determine which nodes within the execution platform 114 are best suited to process the task. Some nodes may have already cached the data needed to process the task (due to the nodes having recently downloaded the data from the cloud computing storage platform 104 for a previous job) and, therefore, be a good candidate for processing the task. Metadata stored in the database 116 assists the compute service manager 112 in determining which nodes in the execution platform 114 have already cached at least a portion of the data needed to process the task. One or more nodes in the execution platform 114 process the task using data cached by the nodes and, if necessary, data retrieved from the cloud computing storage platform 104. It is desirable to retrieve as much data as possible from caches within the execution platform 114 because the retrieval speed is typically much faster than retrieving data from the cloud computing storage platform 104.

[0029] As shown in FIG. 1, the shared data processing platform 100 separates the execution platform 114 from the cloud computing storage platform 104. In this arrangement, the processing resources and cache resources in the execution platform 114 operate independently of the data storage devices 124-1 to 124-N in the cloud computing storage platform 104. Thus, the computing resources and cache resources are not restricted to specific data storage devices 124-1 to 124-N. Instead, all computing resources and all cache resources may retrieve data from, and store data to, any of the data storage resources in the cloud computing storage platform 104.

[0030] FIG. 2 is a block diagram illustrating components of the compute service manager 112, in accordance with some embodiments of the present disclosure. As shown in FIG. 2, a request processing service 202 manages received data storage requests and data retrieval requests (e.g., jobs to

be performed on database data). For example, the request processing service 202 may determine the data necessary to process a received query (e.g., a data storage request or data retrieval request). The data may be stored in a cache within the execution platform 114 or in a data storage device in cloud computing storage platform 104. A management console service 204 supports access to various systems and processes by administrators and other system managers. Additionally, the management console service 204 may receive a request to execute a job and monitor the workload on the system.

[0031] The compute service manager 112 also includes a job compiler 206, a job optimizer 208, and a job executor 210. The job compiler 206 parses a job into multiple discrete tasks and generates the execution code for each of the multiple discrete tasks. The job optimizer 208 determines the best method to execute the multiple discrete tasks based on the data that needs to be processed. The job optimizer 208 also handles various data pruning operations and other data optimization techniques to improve the speed and efficiency of executing the job. The job executor 210 executes the execution code for jobs received from a queue or determined by the compute service manager 112.

[0032] A job scheduler and coordinator 212 sends received jobs to the appropriate services or systems for compilation, optimization, and dispatch to the execution platform 114. For example, jobs may be prioritized and processed in that prioritized order. In an embodiment, the job scheduler and coordinator 212 determines a priority for internal jobs that are scheduled by the compute service manager 112 with other “outside” jobs such as user queries that may be scheduled by other systems in the database but may utilize the same processing resources in the execution platform 114. In some embodiments, the job scheduler and coordinator 212 identifies or assigns particular nodes in the execution platform 114 to process particular tasks. A virtual warehouse manager 214 manages the operation of multiple virtual warehouses implemented in the execution platform 114. As discussed below, each virtual warehouse includes multiple execution nodes that each include a cache and a processor (e.g., a virtual machine, an operating system level container execution environment).

[0033] Additionally, the compute service manager 112 includes a configuration and metadata manager 216, which manages the information related to the data stored in the remote data storage devices and in the local caches (i.e., the caches in execution platform 114). The configuration and metadata manager 216 uses the metadata to determine which data micro-partitions need to be accessed to retrieve data for processing a particular task or job. A monitor and workload analyzer 218 oversees processes performed by the compute service manager 112 and manages the distribution of tasks (e.g., workload) across the virtual warehouses and execution nodes in the execution platform 114. The monitor and workload analyzer 218 also redistributes tasks, as needed, based on changing workloads throughout the network-based database system 102 and may further redistribute tasks based on a user (e.g., “external”) query workload that may also be processed by the execution platform 114. The configuration and metadata manager 216 and the monitor and workload analyzer 218 are coupled to a data storage device 220. Data storage device 220 in FIG. 2 represent any data storage device within the network-based database system 102. For example, data storage device 220 may repre-

sent caches in execution platform **114**, storage devices in cloud computing storage platform **104**, or any other storage device.

[0034] FIG. 3 is a block diagram illustrating components of the execution platform **114**, in accordance with some embodiments of the present disclosure. As shown in FIG. 3, execution platform **114** includes multiple virtual warehouses, which are elastic clusters of compute instances, such as virtual machines. In the example illustrated, the virtual warehouses include virtual warehouse 1, virtual warehouse 2, and virtual warehouse n. Each virtual warehouse (e.g., EC2 cluster) includes multiple execution nodes (e.g., virtual machines) that each include a data cache and a processor. The virtual warehouses can execute multiple tasks in parallel by using the multiple execution nodes. As discussed herein, execution platform **114** can add new virtual warehouses and drop existing virtual warehouses in real time based on the current processing needs of the systems and users. This flexibility allows the execution platform **114** to quickly deploy large amounts of computing resources when needed without being forced to continue paying for those computing resources when they are no longer needed. All virtual warehouses can access data from any data storage device (e.g., any storage device in cloud computing storage platform **104**).

[0035] Although each virtual warehouse shown in FIG. 3 includes three execution nodes, a particular virtual warehouse may include any number of execution nodes. Further, the number of execution nodes in a virtual warehouse is dynamic, such that new execution nodes are created when additional demand is present, and existing execution nodes are deleted when they are no longer necessary (e.g., upon a query or job completion).

[0036] Each virtual warehouse is capable of accessing any of the data storage devices **124-1** to **124-N** shown in FIG. 1. Thus, the virtual warehouses are not necessarily assigned to a specific data storage device **124-1** to **124-N** and, instead, can access data from any of the data storage devices **124-1** to **124-N** within the cloud computing storage platform **104**. Similarly, each of the execution nodes shown in FIG. 3 can access data from any of the data storage devices **124-1** to **124-N**. For instance, the storage device **124-1** of a first user (e.g., provider account user) may be shared with a worker node in a virtual warehouse of another user (e.g., consumer account user), such that the other user can create a database (e.g., read-only database) and use the data in storage device **124-1** directly without needing to copy the data (e.g., copy it to a new disk managed by the consumer account user). In some embodiments, a particular virtual warehouse or a particular execution node may be temporarily assigned to a specific data storage device, but the virtual warehouse or execution node may later access data from any other data storage device.

[0037] In the example of FIG. 3, virtual warehouse 1 includes three execution nodes **302-1**, **302-2**, and **302-N**. Execution node **302-1** includes a cache **304-1** and a processor **306-1**. Execution node **302-2** includes a cache **304-2** and a processor **306-2**. Execution node **302-N** includes a cache **304-N** and a processor **306-N**. Each execution node **302-1**, **302-2**, and **302-N** is associated with processing one or more data storage and/or data retrieval tasks. For example, a virtual warehouse may handle data storage and data retrieval tasks associated with an internal service, such as a clustering service, a materialized view refresh service, a file compac-

tion service, a storage procedure service, or a file upgrade service. In other implementations, a particular virtual warehouse may handle data storage and data retrieval tasks associated with a particular data storage system or a particular category of data.

[0038] Similar to virtual warehouse 1 discussed above, virtual warehouse 2 includes three execution nodes **312-1**, **312-2**, and **312-N**. Execution node **312-1** includes a cache **314-1** and a processor **316-1**. Execution node **312-2** includes a cache **314-2** and a processor **316-2**. Execution node **312-N** includes a cache **314-N** and a processor **316-N**. Additionally, virtual warehouse 3 includes three execution nodes **322-1**, **322-2**, and **322-N**. Execution node **322-1** includes a cache **324-1** and a processor **326-1**. Execution node **322-2** includes a cache **324-2** and a processor **326-2**. Execution node **322-N** includes a cache **324-N** and a processor **326-N**.

[0039] In some embodiments, the execution nodes shown in FIG. 3 are stateless with respect to the data the execution nodes are caching. For example, these execution nodes do not store or otherwise maintain state information about the execution node, or the data being cached by a particular execution node. Thus, in the event of an execution node failure, the failed node can be transparently replaced by another node. Since there is no state information associated with the failed execution node, the new (replacement) execution node can easily replace the failed node without concern for recreating a particular state.

[0040] Although the execution nodes shown in FIG. 3 each include one data cache and one processor, alternative embodiments may include execution nodes containing any number of processors and any number of caches. Additionally, the caches may vary in size among the different execution nodes. The caches shown in FIG. 3 store, in the local execution node (e.g., local disk), data that was retrieved from one or more data storage devices in cloud computing storage platform **104** (e.g., S3 objects recently accessed by the given node). In some example embodiments, the cache stores file headers and individual columns of files as a query downloads only columns necessary for that query.

[0041] To improve cache hits and avoid overlapping redundant data stored in the node caches, the job optimizer **208** assigns input file sets to the nodes using a consistent hashing scheme to hash over table file names of the data accessed (e.g., data in database **116** or database **122**). Subsequent or concurrent queries accessing the same table file will therefore be performed on the same node, according to some example embodiments.

[0042] As discussed, the nodes and virtual warehouses may change dynamically in response to environmental conditions (e.g., disaster scenarios), hardware/software issues (e.g., malfunctions), or administrative changes (e.g., changing from a large cluster to smaller cluster to lower costs). In some example embodiments, when the set of nodes changes, no data is reshuffled immediately. Instead, the least recently used replacement policy is implemented to eventually replace the lost cache contents over multiple jobs. Thus, the caches reduce or eliminate the bottleneck problems occurring in platforms that consistently retrieve data from remote storage systems. Instead of repeatedly accessing data from the remote storage devices, the systems and methods described herein access data from the caches in the execution nodes, which is significantly faster and avoids the bottleneck problem discussed above. In some embodiments,

the caches are implemented using high-speed memory devices that provide fast access to the cached data. Each cache can store data from any of the storage devices in the cloud computing storage platform **104**.

[0043] Further, the cache resources and computing resources may vary between different execution nodes. For example, one execution node may contain significant computing resources and minimal cache resources, making the execution node useful for tasks that require significant computing resources. Another execution node may contain significant cache resources and minimal computing resources, making this execution node useful for tasks that require caching of large amounts of data. Yet another execution node may contain cache resources providing faster input-output operations, useful for tasks that require fast scanning of large amounts of data. In some embodiments, the execution platform **114** implements skew handling to distribute work amongst the cache resources and computing resources associated with a particular execution, where the distribution may be further based on the expected tasks to be performed by the execution nodes. For example, an execution node may be assigned more processing resources if the tasks performed by the execution node become more processor-intensive. Similarly, an execution node may be assigned more cache resources if the tasks performed by the execution node require a larger cache capacity. Further, some nodes may be executing much slower than others due to various issues (e.g., virtualization issues, network overhead). In some example embodiments, the imbalances are addressed at the scan level using a file stealing scheme. In particular, whenever a node process completes scanning its set of input files, it requests additional files from other nodes. If the one of the other nodes receives such a request, the node analyzes its own set (e.g., how many files are left in the input file set when the request is received), and then transfers ownership of one or more of the remaining files for the duration of the current job (e.g., query). The requesting node (e.g., the file stealing node) then receives the data (e.g., header data) and downloads the files from the cloud computing storage platform **104** (e.g., from data storage device **124-1**), and does not download the files from the transferring node. In this way, lagging nodes can transfer files via file stealing in a way that does not worsen the load on the lagging nodes.

[0044] Although virtual warehouses 1, 2, and n are associated with the same execution platform **114**, the virtual warehouses may be implemented using multiple computing systems at multiple geographic locations. For example, virtual warehouse 1 can be implemented by a computing system at a first geographic location, while virtual warehouses 2 and n are implemented by another computing system at a second geographic location. In some embodiments, these different computing systems are cloud-based computing systems maintained by one or more different entities.

[0045] Additionally, each virtual warehouse is shown in FIG. 3 as having multiple execution nodes. The multiple execution nodes associated with each virtual warehouse may be implemented using multiple computing systems at multiple geographic locations. For example, an instance of virtual warehouse 1 implements execution nodes **302-1** and **302-2** on one computing platform at a geographic location and implements execution node **302-N** at a different computing platform at another geographic location. Selecting

particular computing systems to implement an execution node may depend on various factors, such as the level of resources needed for a particular execution node (e.g., processing resource requirements and cache requirements), the resources available at particular computing systems, communication capabilities of networks within a geographic location or between geographic locations, and which computing systems are already implementing other execution nodes in the virtual warehouse.

[0046] Execution platform **114** is also fault tolerant. For example, if one virtual warehouse fails, that virtual warehouse is quickly replaced with a different virtual warehouse at a different geographic location.

[0047] A particular execution platform **114** may include any number of virtual warehouses. Additionally, the number of virtual warehouses in a particular execution platform is dynamic, such that new virtual warehouses are created when additional processing and/or caching resources are needed. Similarly, existing virtual warehouses may be deleted when the resources associated with the virtual warehouse are no longer necessary.

[0048] In some embodiments, the virtual warehouses may operate on the same data in cloud computing storage platform **104**, but each virtual warehouse has its own execution nodes with independent processing and caching resources. This configuration allows requests on different virtual warehouses to be processed independently and with no interference between the requests. This independent processing, combined with the ability to dynamically add and remove virtual warehouses, supports the addition of new processing capacity for new users without impacting the performance observed by the existing users.

[0049] As mentioned above, data from a client storage can be uploaded to the data warehouse. Some techniques can use a “copy” command for this transfer. The “copy” command is typically manually performed or performed based on a set schedule (say, every 15 minutes). However, the use of such “copy” commands can add latency.

[0050] Consequently, latency can be improved by implementing auto-ingestion techniques, as described in further detail below. FIG. 4 is a simplified block diagram of system **400** for automated data ingestion, according to some example embodiments. The system may include a storage **402**, which may be provided as cloud storage (e.g., Amazon S3 storage, Azure storage, GCP storage, etc.). The storage **402** may include client data to upload to the data warehouse.

[0051] The storage **402** may store files (or data) to be ingested into a database **410**. In some embodiments, the storage **402** may include a storage unit **402.1**, an event block **402.2**, and a queue **402.3**. The system may also include a deployment to ingest data in the database **410**. A deployment may include multiple components such as a metadata store/DB, a front-end layer, a load balancing layer, a data warehouse, etc., as discussed above with respect to FIGS. 1-3. The deployments may be provided as public or private deployments. A public deployment may be implemented as a multi-tenant environment, where each tenant or account shares processing and/or storage resources. For example, in a public deployment, multiple accounts may share a metadata store, a front-end layer, a load balancing layer, a data warehouse, etc. A private deployment, on the other hand, may be implemented as a dedicated, isolated environment, where processing and/or storage resources may be dedicated.

[0052] The deployment may be communicatively coupled to the queue 402.3, and may include an integration 404, a pipe 406, and a receiver 408. Integration 404 may be configured to receive a notification when new data becomes available in queue 402.3. For example, the queue may include a pool of Simple Queue Service™ (SQS) queues as part of an Amazon Web Services™ S3 bucket. The pool of SQS queues may be provided to client accounts to add user files to a bucket. A notification may be automatically generated when one or more user files are added to a client account data bucket. A plurality of customer data buckets may be provided for each client account. The automatically generated notification may be received by the integration 404.

[0053] For example, the integration 404 may provide information relating to an occurrence of an event in the queue 402.3. Events may include creation of new data, update of old data, and deletion of old data. The integration 404 may also provide identification information for a resource associated with the event, e.g., the user file that has been created, updated, or deleted. The integration 404 may communicate with the queue 402.3 because the integration 404 may be provided with credentials for the queue 402.3, for example by an administrator and/or user. In an embodiment, the integration 404 may poll the queue 402.3 for notifications.

[0054] The integration 404 may deliver the notification to the pipe 406, which may be provided as a single pipe or multiple pipes. The pipe 406 may store information relating to what data and the location of the data for automatic data ingestion related to the queue 402.3.

[0055] The receiver 408 may perform the automated data ingestion, and then store the ingested data in the database 410. Data ingestion may be performed using the techniques described in U.S. patent application Ser. No. 16/201,854, entitled “Batch Data Ingestion in Database Systems,” filed on Nov. 27, 2018, which is incorporated herein by reference in its entirety, including but not limited to those portions that specifically appear hereinafter, the incorporation by reference being made with the following exception: In the event that any portion of the above-referenced application is inconsistent with this application, this application supercedes the above-referenced application.

[0056] FIG. 5 is a schematic block diagram of a process 500 of ingesting data into a database, according to some example embodiments. The process 500 begins and a storage 502 sends an ingest request, such as a notification. The storage 502 may directly or indirectly communicate with the database system to send in the ingest request. In some embodiments, the ingest request is a notification provided by a third-party vendor storage account, or the ingest request may arise from a compute service manager polling a data lake associated with the client account to determine whether any user files have been added to the client account that have not yet been ingested into the database. The notification includes a list of files to insert into a table of the database. The files are persisted in a queue specific to the receiving table of the database.

[0057] The ingest request is received by a compute service manager 504. The compute service manager 504 identifies at step 506 a user file to ingest. At step 508, the compute service manager identifies a cloud provider type associated with the client account. At step 510, the compute service manager 504 may assign the user file to one or more

execution nodes, based at least in part on the detected cloud provider type, and registers at step 512 micro-partition metadata associated with a database table after the file is ingested into a micro-partition of the database table. The compute service manager 504 provisions one or more execution nodes 516, 520 of an execution platform 514 to perform one or more tasks associated with ingesting the user file. Such ingest tasks 518a, 518b, 522a, 522b include, for example, cutting a file into one or more sections, generating a new micro-partition based on the user file, and/or inserting the new micro-partition in a table of the database.

[0058] The process 500 begins an ingest task that is executed by a warehouse. The ingest task may pull user files from the queue for a database table until it is told to stop doing so. The ingest task may periodically cut a new user file and add it to the database table. In one embodiment, the ingest process is “serverless” in that it is an integrated service provided by the database or compute service manager 504. That is, a user associated with the client account need not provision its own warehouse or a third-party warehouse in order to perform the ingestion process. For example, the database or database provided (e.g., via instances of the compute service manager 504) may maintain the ingest warehouse that then services one or more or all accounts/customers of the database provider.

[0059] In some embodiments, there may be more than one ingest task pulling from a queue for a given table, and this might be necessary to keep up with the rate of incoming data. In some embodiments, the ingest task may decide the time to cut a new file to increase the chances of getting an ideal sized file and avoid “odd sized” files that would result if the file size was lined up with one or more user files. This may come at the cost of added complexity as the track line number of the files consumed must be tracked.

[0060] Users may have files in custom file formats that may not be amenable to ingestion. For example, a data system may support ingestion of native formats, such as csv, json, avro, parquet, orc, xml. However, users may have data in other formats. For example, users may have data in structured file formats, such as syslogs, HL7 messages, VCF, EBDIC, KDF files, dicom images, and complex file formats, such as images, pdf, video, audio, etc., that are not typically supported by data systems. Custom file formats may include formats not natively supported by the data system. Unstructured files (e.g., images) may also be considered custom file formats. Some data systems may have workarounds for custom file formats such as using tasks that have INSERT statements to insert data in the custom file format into a table. However, this workaround suffers from issues. Tasks typically do not have overlapping executions so work will be queued until the previous execution is completed, which adds to latency. Scaling to large volume of continuous events cannot be performed.

[0061] Techniques for continuous ingestion of custom file formats are described below. FIG. 6 is a flow diagram of a method 600 for auto-ingestion of custom file formats, according to some example embodiments. Method 600 may be performed using the auto ingestion techniques described above with reference to FIGS. 4 and 5. That is, a compute service manager, as described herein (e.g., compute service manager 504), may receive notifications of files in a custom file format to be ingested and create a query plan for ingesting those files. The compute service manager may

assign ingest tasks to execution nodes of one or more execution platforms (XPs) as described above.

[0062] At operation **602**, a custom file format may be created for the data system. For example, a “file format” command, which is used to set natively supported file formats, can be extended to include a “custom” option (e.g., “TYPE=CUSTOM”).

[0063] At operation **604**, a user defined table function (UDTF, also referred to as user defined function (UDF)) for the custom file format may be defined. The UDTF can be used to parse the custom file format and return a table. For example, the UDTF can return a table with one or more rows and a variant column. The UDTF may function as a scanner for auto-ingestion, as described herein. The UDTF may include a first argument, which is related to the file to be ingested, and a second argument, which is related to scanner options. The first argument may include a reference to a stage from which the file to be ingested can be read, such as a stage file URL or a scoped URL. The stage may be a cloud storage location (e.g., S3 storage system). From the stage, a stream of the file to be ingested may be obtained.

[0064] The UDTF can parse the stream using parsing code. The parsing code may read bytes from the stream of the file to be ingested. The parsing code may be written in a different programming language, such as java, python, etc., that what is used by the data system. The parsing code may be specific to the custom file format (e.g., HL7). In some examples, the parsing code can use third party libraries (e.g., File Reference) or users can write their own code. The parsing code may decrypt the file and decompress it, if needed. The parsing code returns a stream of row set with at least one variant column.

[0065] The UDTF may be provided in a different programming language (e.g., java, python) that what is used in the data system. The UDTF therefore may be treated as untrusted code by the data system and the UDTF may be executed in a sandbox. In computer security, a sandbox (e.g., sandbox environment) is a security mechanism for separating running programs, usually to prevent system failures or prevent exploitation of software vulnerabilities. A sandbox can be used to execute untested or untrusted programs or code, possibly from unverified or untrusted third parties, suppliers, users or websites, without risking harm to the host machine or operating system. A sandbox can provide a tightly controlled set of resources for guest programs to run in, such as storage and memory scratch space. Network access, the ability to inspect the host system or read from input devices can be disallowed or restricted. UDTFs typically can run in a sandbox environment.

[0066] A sandbox process, in an example, is a program that reduces the risk of security breaches by restricting the running environment of untrusted applications using security mechanisms such as namespaces and secure computing modes (e.g., using a system call filter to an executing process and all its descendants, thus reducing the attack surface of the kernel of a given operating system). Moreover, in an example, the sandbox process is a lightweight process in comparison to an execution node process and is optimized (e.g., closely coupled to security mechanisms of a given operating system kernel) to process a database query in a secure manner within the sandbox environment. In some embodiments, the UDTF is executed using UDTF server, which is restricted from accessing certain files and file systems. For example, the UDTF server and a worker

process handling other operations for the continuous auto-ingestion may be provided as different processors on the same machine.

[0067] In some embodiments, the sandbox process can utilize a virtual network connection in order to communicate with other components within the subject system. A specific set of rules can be configured for the virtual network connection with respect to other components of the subject system. For example, such rules for the virtual network connection can be configured for a particular UDTF to restrict the locations (e.g., particular sites on the Internet or components that the UDTF can communicate) that are accessible by operations performed by the UDTF. Thus, in this example, the UDTF can be denied access to particular network locations or sites on the Internet.

[0068] The sandbox process can be understood as providing a constrained computing environment for a process (or processes) within the sandbox, where these constrained processes can be controlled and restricted to limit access to certain computing resources.

[0069] Examples of security mechanisms can include the implementation of namespaces in which each respective group of processes executing within the sandbox environment has access to respective computing resources (e.g., process IDs, hostnames, user IDs, file names, names associated with network access, and inter-process communication) that are not accessible to another group of processes (which may have access to a different group of resources not accessible by the former group of processes), other container implementations, and the like. By having the sandbox process execute as a sub-process to the execution node process, in some embodiments, latency in processing a given database query can be substantially reduced (e.g., a reduction in latency by a factor of 10x in some instances) in comparison with other techniques that may utilize a virtual machine solution by itself.

[0070] The sandbox process can utilize a sandbox policy to enforce a given security policy. The sandbox policy can be a file with information related to a configuration of the sandbox process and details regarding restrictions, if any, and permissions for accessing and utilizing system resources. Example restrictions can include restrictions to network access, or file system access (e.g., remapping file system to place files in different locations that may not be accessible, other files can be mounted in different locations, and the like). The sandbox process restricts the memory and processor (e.g., CPU) usage of the user code runtime, ensuring that other operations on the same execution node can execute without running out of resources.

[0071] As mentioned above, the sandbox process is a sub-process (or separate process) from the execution node process, which in practice means that the sandbox process resides in a separate memory space than the execution node process. In an occurrence of a security breach in connection with the sandbox process (e.g., by errant or malicious code from a given UDTF), if arbitrary memory is accessed by a malicious actor, the data or information stored by the execution node process is protected.

[0072] At operation **606**, scanner options for the custom file format may be set. Scanner options may relate to how to handle various properties of the data in the file to be ingested, which is in the custom file format. The scanner options may include options related to compression, record delimiter, field delimiter, file extension, skip reader, skipping

blank lines, date format, time format, timestamp format, binary format, null values, column count mismatch.

[0073] At operation 608, pipe definition for one or more pipes may be set to include the custom file format. Moreover, notification of an ingest request for file(s) in the custom file format may be received, as described above. At operation 610, one or more XPs may be pinned for the auto-ingestion of the file(s) to be ingested in the custom file format. As described in further detail below, because the UDTF may include untrusted code, designated or pinned XPs may be used for the auto-ingestion of the custom file format. The pinned one or more XPs may use the pipe(s) set for the custom file format for performing the continuous ingestion.

[0074] At operation 612, the pinned XPs may perform the auto-ingestion of the files to be ingested in the custom file format. The stream output of rows with one or more variant columns of the UDTF may be buffered by the pipe. After the number of buffered rows exceed a threshold value, the pipe may ingest the buffered rows into the source table, as described above. The rows may be created into format files used by the data system and may be registered and committed to the source table. At operation 614, after all files are ingested, the pinned XPs may be released so that they can be used for other tasks, including other ingest tasks.

[0075] FIG. 7 illustrates a data system framework for ingestion of files with custom file formats, according to some example embodiments. The framework may include an ingest compute service manager 702, an ingest warehouse 704, a set of multi-tenant XPs 706.1-706.4, a first set of pinned XPs 708.1-708.2, and a second set of pinned XPs 710.1-710.2. The ingest compute service manager 702 may include the functionalities described above for compute service manager described herein (e.g., compute service manager 112, compute service manager 504). The ingest compute service manager 702 may receive notifications of ingest request from a plurality of accounts associated with the multi-tenant data system. The compute service manager 702 may assign ingest tasks to one or more XPs based on the ingest request notifications. The ingest requests may be for natively supported file formats and/or custom file formats, as described herein.

[0076] The ingest warehouse 704 may include a pool of XPs to perform the ingest tasks assigned by the compute service manager 702. For files in natively supported formats, multi-tenant XPs 706.1-706.4 may be used to perform the ingest tasks. Multi-tenant XPs 706.1-706.4 may perform ingest tasks for multiple accounts at the same time for natively supported format files. However, as described above, for ingest tasks related to custom file formats, pinned XPs are used because the UDTFs defining the custom file format in the pipe definition may include untrusted code. The pinned XPs may be referred to as per-account XP instances for pipes that use custom file formats.

[0077] For example, pinned XPs 708.1-708.2 may be acquired from the pool and pinned by Account A to perform ingest tasks for Account A that are in a custom file format defined by Account A. Because the XPs 708.1-708.2 are pinned while performing ingest tasks for Account A, they cannot perform ingest tasks for other accounts at the same time. After the XPs 708.1-708.2 complete the ingest tasks associated with the ingest request from Account A for files in the custom file format, XPs 708.1-708.2 can be released and be used for other tasks.

[0078] Likewise, pinned XPs 710.1-710.2 may be acquired from the pool and pinned by Account B to perform ingest tasks for Account B that are in a custom file format defined by Account B. Because the XPs 710.1-710.2 are pinned while performing ingest tasks for Account B, they cannot perform ingest tasks for other accounts at the same time. After the XPs 710.1-710.2 complete the ingest tasks associated with the ingest request from Account B for files in a custom file format, XPs 710.1-710.2 can be released and can be used for other tasks. In some examples, smaller side XPs may be used for performing ingest tasks for files with custom file formats so that larger XPs may be reserved for handling multi-tenant ingest tasks.

[0079] FIG. 8 illustrates a diagrammatic representation of a machine 800 in the form of a computer system within which a set of instructions may be executed for causing the machine 800 to perform any one or more of the methodologies discussed herein, according to an example embodiment. Specifically, FIG. 8 shows a diagrammatic representation of the machine 800 in the example form of a computer system, within which instructions 816 (e.g., software, a program, an application, an applet, an app, or other executable code) for causing the machine 800 to perform any one or more of the methodologies discussed herein may be executed. For example, the instructions 816 may cause the machine 800 to execute any one or more operations of any one or more of the methods described herein. As another example, the instructions 816 may cause the machine 800 to implement portions of the data flows described herein. In this way, the instructions 816 transform a general, non-programmed machine into a particular machine 800 (e.g., the remote computing device 106, the access management system 110, the compute service manager 112, the execution platform (XP) 114, the access management system 118, the Web proxy 120) that is specially configured to carry out any one of the described and illustrated functions in the manner described herein.

[0080] In alternative embodiments, the machine 800 operates as a standalone device or may be coupled (e.g., networked) to other machines. In a networked deployment, the machine 800 may operate in the capacity of a server machine or a client machine in a server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine 800 may comprise, but not be limited to, a server computer, a client computer, a personal computer (PC), a tablet computer, a laptop computer, a netbook, a smart phone, a mobile device, a network router, a network switch, a network bridge, or any machine capable of executing the instructions 816, sequentially or otherwise, that specify actions to be taken by the machine 800. Further, while only a single machine 800 is illustrated, the term “machine” shall also be taken to include a collection of machines 800 that individually or jointly execute the instructions 816 to perform any one or more of the methodologies discussed herein.

[0081] The machine 800 includes processors 810, memory 830, and input/output (I/O) components 850 configured to communicate with each other such as via a bus 802. In an example embodiment, the processors 810 (e.g., a central processing unit (CPU), a reduced instruction set computing (RISC) processor, a complex instruction set computing (CISC) processor, a graphics processing unit (GPU), a digital signal processor (DSP), an application-specific integrated circuit (ASIC), a radio-frequency integrated circuit

(RFIC), another processor, or any suitable combination thereof) may include, for example, a processor **812** and a processor **814** that may execute the instructions **816**. The term “processor” is intended to include multi-core processors **810** that may comprise two or more independent processors (sometimes referred to as “cores”) that may execute instructions **816** contemporaneously. Although FIG. **8** shows multiple processors **810**, the machine **800** may include a single processor with a single core, a single processor with multiple cores (e.g., a multi-core processor), multiple processors with a single core, multiple processors with multiple cores, or any combination thereof.

[0082] The memory **830** may include a main memory **832**, a static memory **834**, and a storage unit **836**, all accessible to the processors **810** such as via the bus **802**. The main memory **832**, the static memory **834**, and the storage unit **836** store the instructions **816** embodying any one or more of the methodologies or functions described herein. The instructions **816** may also reside, completely or partially, within the main memory **832**, within the static memory **834**, within the storage unit **836**, within at least one of the processors **810** (e.g., within the processor’s cache memory), or any suitable combination thereof, during execution thereof by the machine **800**.

[0083] The I/O components **850** include components to receive input, provide output, produce output, transmit information, exchange information, capture measurements, and so on. The specific I/O components **850** that are included in a particular machine **800** will depend on the type of machine. For example, portable machines such as mobile phones will likely include a touch input device or other such input mechanisms, while a headless server machine will likely not include such a touch input device. It will be appreciated that the I/O components **850** may include many other components that are not shown in FIG. **8**. The I/O components **850** are grouped according to functionality merely for simplifying the following discussion and the grouping is in no way limiting. In various example embodiments, the I/O components **850** may include output components **852** and input components **854**. The output components **852** may include visual components (e.g., a display such as a plasma display panel (PDP), a light emitting diode (LED) display, a liquid crystal display (LCD), a projector, or a cathode ray tube (CRT)), acoustic components (e.g., speakers), other signal generators, and so forth. The input components **854** may include alphanumeric input components (e.g., a keyboard, a touch screen configured to receive alphanumeric input, a photo-optical keyboard, or other alphanumeric input components), point-based input components (e.g., a mouse, a touchpad, a trackball, a joystick, a motion sensor, or another pointing instrument), tactile input components (e.g., a physical button, a touch screen that provides location and/or force of touches or touch gestures, or other tactile input components), audio input components (e.g., a microphone), and the like.

[0084] Communication may be implemented using a wide variety of technologies. The I/O components **850** may include communication components **864** operable to couple the machine **800** to a network **880** or devices **870** via a coupling **882** and a coupling **872**, respectively. For example, the communication components **864** may include a network interface component or another suitable device to interface with the network **880**. In further examples, the communication components **864** may include wired communication

components, wireless communication components, cellular communication components, and other communication components to provide communication via other modalities. The devices **870** may be another machine or any of a wide variety of peripheral devices (e.g., a peripheral device coupled via a universal serial bus (USB)). For example, as noted above, the machine **800** may correspond to any one of the remote computing device **106**, the access management system **118**, the compute service manager **112**, the execution platform **114**, the Web proxy **120**, and the devices **870** may include any other of these systems and devices.

[0085] The various memories (e.g., **830**, **832**, **834**, and/or memory of the processor(s) **810** and/or the storage unit **836**) may store one or more sets of instructions **816** and data structures (e.g., software) embodying or utilized by any one or more of the methodologies or functions described herein. These instructions **816**, when executed by the processor(s) **810**, cause various operations to implement the disclosed embodiments.

[0086] As used herein, the terms “machine-storage medium,” “device-storage medium,” and “computer-storage medium” mean the same thing and may be used interchangeably in this disclosure. The terms refer to a single or multiple storage devices and/or media (e.g., a centralized or distributed database, and/or associated caches and servers) that store executable instructions and/or data. The terms shall accordingly be taken to include, but not be limited to, solid-state memories, and optical and magnetic media, including memory internal or external to processors. Specific examples of machine-storage media, computer-storage media, and/or device-storage media include non-volatile memory, including by way of example semiconductor memory devices, e.g., erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), field-programmable gate arrays (FPGAs), and flash memory devices; magnetic disks such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks. The terms “machine-storage media,” “computer-storage media,” and “device-storage media” specifically exclude carrier waves, modulated data signals, and other such media, at least some of which are covered under the term “signal medium” discussed below.

[0087] In various example embodiments, one or more portions of the network **880** may be an ad hoc network, an intranet, an extranet, a virtual private network (VPN), a local-area network (LAN), a wireless LAN (WLAN), a wide-area network (WAN), a wireless WAN (WWAN), a metropolitan-area network (MAN), the Internet, a portion of the Internet, a portion of the public switched telephone network (PSTN), a plain old telephone service (POTS) network, a cellular telephone network, a wireless network, a Wi-Fi® network, another type of network, or a combination of two or more such networks. For example, the network **880** or a portion of the network **880** may include a wireless or cellular network, and the coupling **882** may be a Code Division Multiple Access (CDMA) connection, a Global System for Mobile communications (GSM) connection, or another type of cellular or wireless coupling. In this example, the coupling **882** may implement any of a variety of types of data transfer technology, such as Single Carrier Radio Transmission Technology (1xRTT), Evolution-Data Optimized (EVDO) technology, General Packet Radio Service (GPRS) technology, Enhanced Data rates for GSM

Evolution (EDGE) technology, third Generation Partnership Project (3GPP) including 3G, fourth generation wireless (4G) networks, Universal Mobile Telecommunications System (UMTS), High-Speed Packet Access (HSPA), Worldwide Interoperability for Microwave Access (WiMAX), Long Term Evolution (LTE) standard, others defined by various standard-setting organizations, other long-range protocols, or other data transfer technology.

[0088] The instructions **816** may be transmitted or received over the network **880** using a transmission medium via a network interface device (e.g., a network interface component included in the communication components **864**) and utilizing any one of a number of well-known transfer protocols (e.g., hypertext transfer protocol (HTTP)). Similarly, the instructions **816** may be transmitted or received using a transmission medium via the coupling **872** (e.g., a peer-to-peer coupling) to the devices **870**. The terms “transmission medium” and “signal medium” mean the same thing and may be used interchangeably in this disclosure. The terms “transmission medium” and “signal medium” shall be taken to include any intangible medium that is capable of storing, encoding, or carrying the instructions **816** for execution by the machine **800**, and include digital or analog communications signals or other intangible media to facilitate communication of such software. Hence, the terms “transmission medium” and “signal medium” shall be taken to include any form of modulated data signal, carrier wave, and so forth. The term “modulated data signal” means a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal.

[0089] The terms “machine-readable medium,” “computer-readable medium,” and “device-readable medium” mean the same thing and may be used interchangeably in this disclosure. The terms are defined to include both machine-storage media and transmission media. Thus, the terms include both storage devices/media and carrier waves/modulated data signals.

[0090] The various operations of example methods described herein may be performed, at least partially, by one or more processors that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Similarly, the methods described herein may be at least partially processor-implemented. For example, at least some of the operations of the methods described herein may be performed by one or more processors. The performance of certain of the operations may be distributed among the one or more processors, not only residing within a single machine, but also deployed across a number of machines. In some example embodiments, the processor or processors may be located in a single location (e.g., within a home environment, an office environment, or a server farm), while in other embodiments the processors may be distributed across a number of locations.

[0091] Although the embodiments of the present disclosure have been described with reference to specific example embodiments, it will be evident that various modifications and changes may be made to these embodiments without departing from the broader scope of the inventive subject matter. Accordingly, the specification and drawings are to be regarded in an illustrative rather than a restrictive sense. The accompanying drawings that form a part hereof show, by way of illustration, and not of limitation, specific embodiments in which the subject matter may be practiced. The

embodiments illustrated are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed herein. Other embodiments may be used and derived therefrom, such that structural and logical substitutions and changes may be made without departing from the scope of this disclosure. This Detailed Description, therefore, is not to be taken in a limiting sense, and the scope of various embodiments is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

[0092] Such embodiments of the inventive subject matter may be referred to herein, individually and/or collectively, by the term “invention” merely for convenience and without intending to voluntarily limit the scope of this application to any single invention or inventive concept if more than one is in fact disclosed. Thus, although specific embodiments have been illustrated and described herein, it should be appreciated that any arrangement calculated to achieve the same purpose may be substituted for the specific embodiments shown. This disclosure is intended to cover any and all adaptations or variations of various embodiments. Combinations of the above embodiments, and other embodiments not specifically described herein, will be apparent, to those of skill in the art, upon reviewing the above description.

[0093] In this document, the terms “a” or “an” are used, as is common in patent documents, to include one or more than one, independent of any other instances or usages of “at least one” or “one or more.” In this document, the term “or” is used to refer to a nonexclusive or, such that “A or B” includes “A but not B,” “B but not A,” and “A and B,” unless otherwise indicated. In the appended claims, the terms “including” and “in which” are used as the plain-English equivalents of the respective terms “comprising” and “wherein.” Also, in the following claims, the terms “including” and “comprising” are open-ended; that is, a system, device, article, or process that includes elements in addition to those listed after such a term in a claim is still deemed to fall within the scope of that claim.

[0094] Described implementations of the subject matter can include one or more features, alone or in combination as illustrated below by way of example.

[0095] Example 1. A method comprising: defining, by a data system comprising at least one hardware processor, a custom file format for use in auto-ingestion of files in the custom file format; setting a user defined function (UDF) to parse files in the custom file format and return a table with one or more rows and a variant column; receiving a notification of at least one file in the custom file format for auto-ingestion; performing auto-ingestion of the at least one file using a pipe set for the custom file format executing the UDF; and storing data from the at least one file in a source table.

[0096] Example 2. The method of example 1, wherein the custom file format is not natively supported by the data system.

[0097] Example 3. The method of any of examples 1-2, wherein the custom file format includes an unstructured file type.

[0098] Example 4. The method of any of examples 1-3, wherein the data system operates using a first programming language and the UDF is in a second programming language.

[0099] Example 5. The method of any of examples 1-4, wherein the UDF is executed in a sandbox environment.

[0100] Example 6. The method of any of examples 1-5, further comprising: acquiring one or more execution nodes; assigning one or more ingest tasks related to the at least one file; pinning the one or more execution nodes such that the one or more execution nodes cannot perform other ingestion tasks for files in a different format than the custom file format; and performing, by the one or more execution nodes, the one or more ingest tasks.

[0101] Example 7. The method of any of examples 1-6, further comprising: releasing the one or more execution nodes such that they can perform other ingestions tasks for files in different formats than the custom file format.

[0102] Example 8. A system comprising: one or more processors of a machine; and a memory storing instructions that, when executed by the one or more processors, cause the machine to perform operations implementing any one of example methods 1 to 7.

[0103] Example 9. A machine-readable storage device embodying instructions that, when executed by a machine, cause the machine to perform operations implementing any one of example methods 1 to 7.

What is claimed is:

1. A method for managing a pool of execution nodes in a multi-tenant data system the method comprising:

in response to identifying a first file of a natively supported format for ingestion, acquiring a first set of execution nodes from the pool of execution nodes;

performing, using the first set of execution nodes, ingestion of the first file, wherein the first set of execution nodes remains available for concurrent assignment of tasks for other accounts of the multi-tenant data system;

in response to identifying a second file of a custom format not being natively supported by the multi-tenant data system for ingestion, acquiring a second set of execution nodes from the pool of execution nodes; and

performing, using the second set of execution nodes, ingestion of the second file using a user-defined function (UDF) to parse the custom format, wherein the second set of execution nodes is pinned during said ingestion of the second file to prevent concurrent assignment of tasks for other accounts of the multi-tenant data system.

2. The method of claim 1, further comprising:

storing data from the second file in one or more rows in a source table in a format used by the multi-tenant data system;

registering the one or more rows in the source table; and committing the one or more rows in the source table.

3. The method of claim 2, further comprising:

in response to committing the one or more rows, releasing the second set of execution nodes.

4. The method of claim 1, wherein parsing the second file returns a table with one or more rows and a variant column.

5. The method of claim 1, wherein the custom format includes an unstructured file type.

6. The method of claim 1, wherein the multi-tenant data system operates using a first programming language and the UDF is in a second programming language.

7. The method of claim 6, wherein the UDF is executed in a sandbox environment.

8. A system comprising:

at least one hardware processor; and

at least one memory storing instructions that, when executed by the at least one hardware processor, cause the at least one hardware processor to perform operations comprising:

in response to identifying a first file of a natively supported format for ingestion, acquiring a first set of execution nodes from a pool of execution nodes in a multi-tenant data system;

performing, using the first set of execution nodes, ingestion of the first file, wherein the first set of execution nodes remains available for concurrent assignment of tasks for other accounts of the multi-tenant data system;

in response to identifying a second file of a custom format not being natively supported by the multi-tenant data system for ingestion, acquiring a second set of execution nodes from the pool of execution nodes; and

performing, using the second set of execution nodes, ingestion of the second file using a user-defined function (UDF) to parse the custom format, wherein the second set of execution nodes is pinned during said ingestion of the second file to prevent concurrent assignment of tasks for other accounts of the multi-tenant data system.

9. The system of claim 8, further comprising:

storing data from the second file in one or more rows in a source table in a format used by the multi-tenant data system;

registering the one or more rows in the source table; and committing the one or more rows in the source table.

10. The system of claim 9, further comprising:

in response to committing the one or more rows, releasing the second set of execution nodes.

11. The system of claim 8, wherein parsing the second file returns a table with one or more rows and a variant column.

12. The system of claim 8, wherein the custom format includes an unstructured file type.

13. The system of claim 8, wherein the multi-tenant data system operates using a first programming language and the UDF is in a second programming language.

14. The system of claim 13, wherein the UDF is executed in a sandbox environment.

15. A machine-storage medium embodying instructions that, when executed by a machine, cause the machine to perform operations comprising:

in response to identifying a first file of a natively supported format for ingestion, acquiring a first set of execution nodes from a pool of execution nodes in a multi-tenant data system;

performing, using the first set of execution nodes, ingestion of the first file, wherein the first set of execution nodes remains available for concurrent assignment of tasks for other accounts of the multi-tenant data system;

in response to identifying a second file of a custom format not being natively supported by the multi-tenant data system for ingestion, acquiring a second set of execution nodes from the pool of execution nodes; and

performing, using the second set of execution nodes, ingestion of the second file using a user-defined function (UDF) to parse the custom format, wherein the second set of execution nodes is pinned during said

ingestion of the second file to prevent concurrent assignment of tasks for other accounts of the multi-tenant data system.

16. The machine-storage medium of claim **15**, further comprising:

storing data from the second file in one or more rows in a source table in a format used by the multi-tenant data system;

registering the one or more rows in the source table; and committing the one or more rows in the source table.

17. The machine-storage medium of claim **16**, further comprising:

in response to committing the one or more rows, releasing the second set of execution nodes.

18. The machine-storage medium of claim **15**, wherein parsing the second file returns a table with one or more rows and a variant column.

19. The machine-storage medium of claim **15**, wherein the custom format includes an unstructured file type.

20. The machine-storage medium of claim **15**, wherein the multi-tenant data system operates using a first programming language and the UDF is in a second programming language, wherein the UDF is executed in a sandbox environment.

* * * * *