

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent Application Publication

20250258811

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

AL-SHIHABI; Noah Amin et al.

---

### **SCHEMA SNAPSHOT ISOLATION ACCESS IN A RELATIONAL DATABASE**

---

#### **Abstract**

Relational database systems are disclosed that are enabled to operate with versioned metadata. The relational database system includes a lock manager, a transaction manager and a version aware metadata storage and cache configured to store to store and manage versions of metadata, to determine which of such versions should be visible at any given point in time, and to enable creation of the proper versions of metadata. In an aspect, the transaction manager manages transaction identifiers and their associated start times, abort times and/or commit times. Such data enables determination of transaction visibility, and consequently the metadata version visibility, for any point in time. In an aspect, such metadata versioning support enables snapshot isolation of metadata transactions.

---

**Inventors:** AL-SHIHABI; Noah Amin (Sattle, WA), ANTONOPOULOS; Panagiotis (Redmond, WA), KODAVALLA; Hanumantha Rao (Sammamish, WA), OGG; Nikolas Christopher (Seattle, WA), TIKMANI; Divyesh Rameshchandra (Hyderabad, Telangana, IN), KIM; Wonseok (Sammamish, WA), SWANSON; Alexander Wai-man (Redmond, WA), JAIN; Rajat (Redmond, WA)

**Applicant:** Microsoft Technology Licensing, LLC (Redmond, WA)

**Family ID:** 95657651

**Appl. No.:** 19/194515

**Filed:** April 30, 2025

#### **Related U.S. Application Data**

parent US continuation 18428887 20240131 PENDING child US 19194515

us-provisional-application US 63598826 20231114

---

## Publication Classification

Int. Cl.: G06F16/23 (20190101); G06F16/2453 (20190101)

U.S. Cl.:

CPC G06F16/2379 (20190101); G06F16/2343 (20190101); G06F16/24542 (20190101);

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 18/428,887, titled “SCHEMA SNAPSHOT ISOLATION ACCESS IN A RELATIONAL DATABASE,” filed Jan. 31, 2024, now allowed, which claims priority to U.S. Provisional Patent Application Ser. No. 63/598,826, filed Nov. 14, 2023, and titled “SCHEMA SNAPSHOT ISOLATION ACCESS IN A RELATIONAL DATABASE,” the entirety of which is incorporated by reference herein.

### BACKGROUND

[0002] Relational database systems (e.g., Microsoft SQL Server) typically require full synchronization between read queries and data definition language operations (“DDLs”) that modify schemas because metadata is not versioned. Thus, all DDLs require a full exclusive lock (i.e., a schema modification lock denoted as SCH-M) on all schema objects being modified to synchronize with any concurrent read queries that require access to these objects. This requirement imposes certain operational limitations.

[0003] For example, a requested SCH-M lock required for a DDL operation may not be granted until all queries that depend on the schema being modified have completed thus blocking the DDL operation and causing a period of unavailability that significantly impacts the execution performance of a workload. Likewise, while the DDL operation is waiting for ongoing queries to complete, new queries submitted after the SCH-M lock is requested are also blocked, further extending the window of unavailability.

[0004] Further, DDL operations invalidate any snapshot transaction that attempts to access the schema objects after the DDL operation commits since the schema associated with the transaction snapshot no longer exists after the DDL operation commits, and thus can no longer be retrieved. This causes unexpected failures of workload queries, and such queries must be retried further impacting performance of the workload.

[0005] Likewise, the fact that DDL operations need to wait for all ongoing queries to complete before they can get the schema modification lock and proceed significantly impacts the health of secondary replicas. Secondary replicas rely on log replication, and such log replication must wait for extended periods of time as it might be blocked behind long-running analytical queries which are common on secondary replicas. This can lead to accumulating a large log on a primary replica and eventually running out of log space. Furthermore, replication of such large logs results in long recovery times and unavailability of the secondary replica when it restarts or when a database is failed over to the secondary replica since it needs to redo all of the accumulated logs for the replica.

### SUMMARY

[0006] This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

[0007] Relational database systems are described herein that are enabled to operate with versioned

metadata. The relational database system includes a lock manager, a transaction manager and a version aware metadata storage and cache configured to store and manage versions of metadata, to determine which of such versions should be visible at any given point in time, and to enable creation of the proper versions of metadata. In an aspect, the transaction manager manages transaction identifiers and their associated start times, abort times and/or commit times. Such data enables determination of transaction visibility, and consequently the metadata version visibility, for any point in time. In an aspect, metadata versioning enables snapshot isolation of metadata transactions.

[0008] Further features and advantages, as well as the structure and operation of various examples, are described in detail below with reference to the accompanying drawings. It is noted that the ideas and techniques are not limited to the specific examples described herein. Such examples are presented herein for illustrative purposes only. Additional examples will be apparent to persons skilled in the relevant art(s) based on the teachings contained herein.

---

## **Description**

### **BRIEF DESCRIPTION OF THE DRAWINGS/FIGURES**

[0009] The accompanying drawings, which are incorporated herein and form a part of the specification, illustrate embodiments of the present application and, together with the description, further serve to explain the principles of the embodiments and to enable a person skilled in the pertinent art to make and use the embodiments.

[0010] FIG. 1 depicts an example relational database system providing versioned metadata, according to an example embodiment.

[0011] FIG. 2 depicts an example logical employee table stored in an example SQL database, according to an example embodiment.

[0012] FIG. 3 depicts an example logical timeline of concurrent data read and metadata modification transactions using snapshot isolation, according to an example embodiment.

[0013] FIG. 4 depicts an example logical timeline of concurrent data read and metadata modification transactions using read committed snapshot isolation, according to an example embodiment.

[0014] FIG. 5A depicts a flowchart of an example method of operating a relational database system to modify a metadata object without being blocked by a concurrent read access of the metadata object, according to an example embodiment.

[0015] FIG. 5B depicts a flowchart of a refinement to the method depicted in FIG. 5A, according to an embodiment.

[0016] FIG. 5C depicts a flowchart of a refinement to the method depicted in FIGS. 5A and 5B, according to an embodiment.

[0017] FIG. 6 depicts a flowchart of a refinement to the methods depicted in FIGS. 5A, 5B and 5C for triggering query plan recompilation in a snapshot isolation mode, according to an embodiment.

[0018] FIG. 7 is a block diagram of an example computer system in which embodiments may be implemented.

[0019] The features and advantages of embodiments will become more apparent from the detailed description set forth below when taken in conjunction with the drawings, in which like reference characters identify corresponding elements throughout. In the drawings, like reference numbers generally indicate identical, functionally similar, and/or structurally similar elements. The drawing in which an element first appears is indicated by the leftmost digit(s) in the corresponding reference number.

### **DETAILED DESCRIPTION**

#### **I. Introduction**

[0020] The following detailed description discloses numerous example embodiments. The scope of the present patent application is not limited to the disclosed embodiments, but also encompasses combinations of the disclosed embodiments, as well as modifications to the disclosed embodiments. It is noted that any section/subsection headings provided herein are not intended to be limiting. Embodiments are described throughout this document, and any type of embodiment may be included under any section/subsection. Furthermore, embodiments disclosed in any section/subsection may be combined with any other embodiments described in the same section/subsection and/or a different section/subsection in any manner.

## II. Example Embodiments

[0021] Embodiments herein are described in the context of a relational database system that provides versioned metadata. For example, Microsoft® SQL Server and similar relational database systems have historically operated as a “single master” database application that lacks metadata versioning and requires metadata modification to be performed by a single process that has full and exclusive access to the metadata objects being modified. These limitations result in severe scalability problems because transaction ACID (Atomicity, Consistency, Isolation and Durability) compliance is achieved at the expense of concurrency. Embodiments disclosed herein, on the other hand, overcome these “single master” metadata issues by providing relational database systems with native metadata versioning and new lock schemes that enable modification of metadata objects concurrently with read operations that reference or depend on such metadata objects. As described further below, the disclosed embodiments of a relational database system using versioned metadata enables, among other things, snapshot isolation of data definition language (“DDL”) operations that modify metadata objects (e.g., a database schema).

[0022] For example, consider FIG. 2 which depicts an example logical employees table **202** that may be stored in an example SQL database, according to an example embodiment. Table **202** includes rows and columns. Each row corresponds to a particular employee whereas each column corresponds to information about each respective employee. Each row may be thought of as a record, and each of the corresponding column entries for a record may be considered data fields of the corresponding record. For example, the first column labeled emp\_id includes the employee ID for each employee. Likewise, other columns include data such as the name, job title and the like. Every table has a corresponding metadata object stored in the database. The metadata object includes, as the name implies, data about data. In particular, data about the data stored in the table to which the object corresponds. Thus, for example, the metadata that corresponds to Table **202** identifies each column and specifies its data type. When, for example, a DDL command causes a column to be added or dropped from a table, the metadata object that corresponds to that table must likewise be modified to reflect the new logical structure of the table.

[0023] Snapshot isolation of metadata transactions enables virtually the same benchmark performance to be maintained by relational database systems with versioned metadata as for relational database systems without versioned metadata. Furthermore, overall system performance is significantly improved by reducing the blocking ordinarily caused by DDL operations (schema changes) that typically issue a great many read operations.

[0024] A general description of example operating environments is provided below, followed by a detailed description of these and further example embodiments.

[0025] Conventional relational database systems require a full exclusive lock to execute metadata modification commands such as DDL statements. The consequence of this is that a snapshot transaction (e.g., a read query) that attempts to access a metadata object after a DDL transaction commits will fail since the original schema based on the transaction snapshot can no longer be retrieved. Embodiments, on the other hand, enable such read transactions to succeed by referencing the version of metadata that existed before the DDL transaction committed.

[0026] In particular, embodiments enable concurrent access to metadata objects while they are being modified by supporting multiple versions of those objects. Each schema modification of an

object (e.g., via DDL command) creates a new version of the object, and other accessors can still access the version of the object that corresponds to the snapshot of their transaction. For many common DDL operations, all this occurs without blocking between common DDL operations and other accesses (i.e., providing full read access to the data during such operations). Similarly, continuous redo is not blocked by long-running analytical queries, so the logs on primaries are enabled to be cleared before getting too large, and recovery time is reduced as a result of a smaller redo log.

[0027] Embodiments include a reconfiguration of metadata layers to support multi-versioning. This is similar to how Multi-Versioning Concurrency Control (MVCC) works for data but in this case, implemented for metadata. The content of system tables and metadata caches is versioned and enables versioned access following the Snapshot Isolation semantics. In examples, introduced algorithms enable:

[0028] 1) Querying for metadata objects based on different attributes (e.g., identifier (ID), name, etc.) and finding the most relevant version of an object even if multiple of these attributes might have changed over time (e.g., when a table is renamed).

[0029] 2) Reloading the cache by discovering all possibly visible versions of a record within the corresponding system tables to guarantee that all the necessary versions of the objects are correctly loaded.

[0030] 3) Versioned accesses of the caches and underlying system tables using the corresponding transaction or statement snapshot information and following the semantics (snapshot isolation modes) of Snapshot Isolation or RCSI (Read Committed Snapshot Isolation).

[0031] 4) Reconfigure how snapshots are established and used in the system to guarantee that query compilation (that accesses metadata) can perform versioned access to load the correct version of the object and not block behind any ongoing DDLs. This algorithm also allows for internal retries to minimize failures that can occur due to schema inconsistencies without surfacing them to the users.

[0032] 5) Reconfigure SQL locking semantics around metadata access by introducing two new lock types that allow us to control which DDL operations and accesses (e.g., depending on the isolation level) should support concurrency or not. Beyond allowing concurrent access between DDLs and queries, our design can support versioned access to schema information for temporal (AS OF) queries. This enables users to access their tables in the same shape/configuration they had at the specified time.

[0033] Embodiments enable advantages over conventional techniques. For instance, embodiments enable advantages over the following conventional techniques:

[0034] 1) Semi-structured data stores, like document stores, that allow for flexible schemas: These systems fundamentally do not face similar challenges due to their loose schema definition that allows different rows/documents to have different schema. Despite this benefit, such systems are challenging for developers to keep their applications and data aligned since it is very easy for the schema to diverge over time. Embodiments disclosed herein enable concurrent DDL and queries while still supporting a strict schema definition.

[0035] 2) RDBMSs (relational database management systems) that support similar functionality. One key difference is that such conventional systems were designed from scratch to support versioned metadata. Embodiments disclosed herein enable retrofitting versioning to existing, non-versioned RDBMSs.

[0036] These and further features, configurations, and advantages are described as follows with reference to FIG. 1. FIG. 1 depicts an example relational database system **100** providing versioned metadata. As shown in FIG. 1, relational database system **100** includes an execution environment **108**, a client **104**, a client **106**. Execution environment **108** includes a binding/algebraizer/query optimizer **112**, a parser **110**, a lock manager **124**, a metadata storage **114**, a data storage engine **118**, a query execution engine **120** and a transaction manager **122**. Metadata storage **114** includes a

metadata cache **116**. Other structural and operational embodiments will be apparent to persons skilled in the relevant art(s) based on the following discussion regarding relational database system **100** as depicted in FIG. **1**.

[0037] The interconnections between each of client **104**, client **106** and execution environment **108** as depicted in FIG. **1** may represent network interconnections, in an embodiment. For example, such network connections are made via the Internet. However, portions of the depicted connections may utilize other networks such as local area networks (LANs), wide area networks (WANs), enterprise networks, and may include one or more of wired and/or wireless portions. Moreover, the interconnections depicted in FIG. **1** may represent direct connections between some or all of the various components. In other embodiments, some or all of the depicted components may be incorporated into a single system whereby the depicted interconnections are logical rather than physical. For example, some or all of transaction manager **122**, data storage engine **118**, query execution engine **120**, and the like, may together or alone be incorporated into the same component as separate logical operating units with the depicted interconnections comprising data flows between the operating units, but internal to the same physical component.

[0038] Clients **104** and **106** in different embodiments are any number, type, or combination of computing devices or computing systems, including a terminal, a personal computer, a laptop computer, a tablet device, a smart phone, a personal digital assistant, a server(s), a gaming console, and/or the like, including internal/external storage devices, that are utilized to execute functions/operations described herein for performing operations on relational database system **100**, e.g., providing queries to a SQL server, as well as for performing client-side functions/operations of client-server scenarios. Clients **104** and **106** also include additional components (not shown for brevity and illustrative clarity) including, but not limited to, components and subcomponents of other devices and/or systems herein, in various embodiments.

[0039] Clients **104** and **106** may be a computing device associated with a domain which, as used herein, generally refers to a physical and/or logical system boundary under the control of an entity within which applications and/or services are hosted, offered, managed, and/or otherwise implemented, and also encompasses subdomains and/or the like in embodiments. Exemplary, non-limiting domains include, without limitation, web domains, tenancies of hosted cloud platforms, cloud service providers, enterprise systems, and/or any other type of network or system. A tenant is particular type of domain that is a representation of an organization in a cloud platform. The domain of the tenant in the cloud platform is its tenancy in which the tenant registers and manages applications, stores data/files, accesses services, etc.

[0040] Data storage engine **118** and metadata storage **114** may comprise or include any type and/or number of data storage devices or systems, and may comprise internal and/or external storage in various embodiments. That is, for example, metadata storage **114** may comprise a monolithic storage device/system, a cloud-based storage system, a distributed storage system, and/or the like.

[0041] Moreover, although each of clients **104** and **106**, metadata storage **114**, lock manager **124**, data storage engine **118**, query execution engine **120** and transaction manager **122** in FIG. **1** are depicted as monolithic components, it should be understood that such components may be embodied in any number of computing devices including servers, and may include any type and number of other resources, including resources that facilitate communications with and between computing devices connected via one of the aforementioned networks. In embodiments, servers implementing these components may be organized in any manner, including being grouped in server racks (e.g., 8-40 servers per rack, referred to as nodes or “blade servers”), server clusters (e.g., 2-64 servers, 4-8 racks, etc.), or datacenters (e.g., thousands of servers, hundreds of racks, dozens of clusters, etc.). In an embodiment, the servers of may be co-located (e.g., housed in one or more nearby buildings with associated components such as backup power supplies, redundant data communications, environmental controls, etc.) to form a datacenter, or may be arranged in other manners. Accordingly, in an embodiment, clients **104** and **106**, metadata storage **114**, lock manager

**124**, data storage engine **118**, query execution engine **120** and transaction manager **122** may comprise a datacenter in a distributed collection of datacenters.

[0042] Each of client **104** and **106** work off the same user database(s), performs updates thereto, and can access data and/or metadata as of any prior point in time, in an embodiment. Moreover, and as will be described further herein below, a schema of the objects in any database may be altered by one of client **104** or **106** while the other client or other clients (not shown) are running a concurrent query that is scanning the same schema. Such alterations and deletions are typically performed by executing DDL operations. For example, and as will be described below, a user of client **104** or **106** may connect to execution environment **108** of relational database system **100** and execute a T-SQL DDL statement. Relational database system **100** may perform the alteration/deletion by executing the command in a transaction. As mentioned above, such transactions must be ACID compliant as will now be described in further detail.

[0043] As mentioned herein above, data and/or metadata modifications performed by relational database system **100** are performed in the context of an ACID compliant transaction, and embodiments of relational database system **100** are ACID compliant. As known in the art, ACID is an acronym that denotes a set of properties of database transactions that ensure that data (or metadata) persisted to the database are valid despite errors that may arise due to, for example, hardware failures or power failures. The ACID properties are: Atomicity, Consistency, Isolation and Durability. Transactions performed by relational database system **100** are ACID compliant inasmuch as the operations that logically correspond to the transaction collectively satisfy the ACID properties.

[0044] The Atomicity property of a transaction demands that the transaction either completely succeed, or completely fail. For a transaction to completely fail means that the database is left unchanged. For example, suppose a transaction comprises transferring money from account A to account B. The entire transaction comprises a number of steps as follows: funds are debited from account A, funds are transferred to wherever account B is held, and funds are credited to account B. In this situation, atomicity guarantees that no funds are debited from account A if, for whatever reason, funds are not credited to account B.

[0045] The Consistency property ensures that a transaction conform to all applicable rules governing the storage of the data such that the transaction brings the database from one valid state to another valid state.

[0046] The Isolation property ensures that different transactions that execute concurrently leave the database in the same state that would have been obtained if the transactions were executed in series. Isolation is the main goal of concurrency control, and embodiments discussed further herein below enforce such concurrency controls on metadata operations with relational database system **100**.

[0047] Finally, the Durability property guarantees that once a transaction is committed (i.e., completed and persisted to the database in an ACID compliant manner), the transaction will remain committed and that no hardware, system or power failure can cause the transaction to be lost or for the database to enter an otherwise inconsistent state. With further reference to relational database system **100** of FIG. 1, ACID properties of transactions performed by relational database system **100** are ensured, in part, through the use of logging performed as follows.

[0048] In order to maintain the ACID properties of the transaction, embodiments of relational database system **100** may be configured to generate a log record for a transaction when the transaction commits and store that record in a transaction log before any data modifications caused by the transaction are written to disk. In an embodiment, logging may be handled by, for example, transaction manager **122** or data storage engine **118**, alone or in combination with one another. In other embodiments, however, a dedicated logging component (not shown in FIG. 1) may be employed to manage the storage and retrieval of such transaction logs.

[0049] A log record for a committed transaction includes all the information necessary to re-do the

transaction in the event there is a problem (e.g., power failure) before the data modified by the transaction can be stored (e.g., in data storage engine **118**). A log record may comprise information that includes, but is not limited to, a transaction identifier, a log sequence number, a time stamp, information indicating what data object or objects was/were modified and how, and the like. In an embodiment, the log record information is managed and/or provided by transaction manager **122**. [0050] In an embodiment, transaction manager **122** may be configured to maintain a transaction ID and start time for a transaction. Transaction manager **122** may be centralized and available to provide transaction services to all clients and/or replicas of relational database system **100**. Embodiments may be configured to acquire a list of active and aborted transactions as of any point in time from transaction manager **122**, and to use such information for versioning purposes as described further herein below. Furthermore, embodiments may acquire locks through lock manager **124** as needed. Such locks will now be further described.

[0051] Relational database systems use locks to maintain the database in a consistent state by ensuring that any read, scan or other use of metadata objects will reflect a valid state of such metadata objects. Conventional relational database systems do not provide versioned metadata and as such, there can only be one valid state for a metadata object. The consequence of this for conventional relational database systems, and as described above, is that metadata objects that are currently being scanned or otherwise read cannot be modified and likewise, metadata objects that are in the process of being modified cannot be scanned or otherwise read. Locks are the mechanism by which these restrictions are enforced.

[0052] Suppose, for example, that a metadata modification command (i.e., DDL) is being executed in a transaction. To enforce the restrictions of conventional relational database systems described immediately herein above, execution of the DDL will require a lock on the metadata object(s) being modified. Prior to such modification, the relational database system will request an exclusive schema modification lock (denoted as “SCH-M”) from, for example, lock manager **124** or an analogous component. Lock manager **124** will provide and enforce such a lock only after ensuring that no other transaction is reading or scanning the metadata object to be modified. Likewise, a read/scan of a metadata object also requires a lock which similarly prevents a metadata modification from taking place. Such a lock is referred to as a schema stability (SCH-S) lock.

[0053] As a result, the SCH-M and SCH-S locks are mutually exclusive. That is, a lock manager such as, for example, lock manager **124** will not issue a SCH-M lock for a schema change DDL until no SCH-S locks are outstanding. Likewise, a SCH-S lock will not be provided until there are no SCH-M locks outstanding. The overall result for conventional systems is that metadata reads and writes are mutually exclusive and may not be executed concurrently. The disclosed embodiments, on the other hand, allow metadata reads and writes to be executed concurrently through the use of versioned metadata, and by introducing new lock types.

[0054] Versioning-based snapshot isolation of metadata transactions is enabled in part by maintaining transaction versions for every metadata version metadata storage **114** and in metadata cache **116** of metadata storage **114**. Metadata cache **116** provides a high speed access to metadata objects and their various versions, whereas metadata storage **114** includes a physical storage in a non-volatile memory. Of course, the contents of high speed metadata cache **116** must be kept in sync with the physical record stored elsewhere within metadata storage **114**.

[0055] Versioning-based snapshot isolation of metadata transactions is further enabled by providing the means to determine which version of a metadata object is visible to a given transaction. In an embodiment, transaction manager **122** may maintain a transaction visibility list for each transaction that dictates which transactions, and consequently which metadata version(s), are visible to that transaction. In other embodiments, however, some combination of transaction manager **122**, metadata storage **114**, data storage engine **118** and/or lock manager **124** may be configured alone or in combination to enable embodiments to determine which versions of metadata objects are available to a particular transaction. Versioning-based snapshot isolation is thereby enabled in the



following general manner.

[0056] The versioning framework of the disclosed embodiments invoke a copy-on-write mechanism when metadata is modified or deleted. This requires that while the transaction performing such modification or deletion is still running, the old version of the metadata must be available for transactions that require an earlier transactionally consistent state. Accordingly, when a versioning-based isolation level is enabled, metadata storage **114** stores a version of a metadata object each time that object is modified. Each new version of the metadata object may, for example, be stored along with the transaction identifier of the transaction that committed the modification along with a commit timestamp thereby allowing embodiments to determine the point in time when a given version came into existence. Alternatively, committed transactions IDs and associated timestamps may be maintained separately from the versions that correspond to same.

[0057] FIGS. **3** and **4** depict example logical timelines of concurrent metadata read and write transactions under two different snapshot isolation schemes, according to an example embodiment. In particular, FIG. **3** depicts an example logical timeline **300** of concurrent data read and metadata modification transactions using example snapshot isolation **314**, whereas FIG. **4** depicts an example logical timeline **400** of concurrent data read and metadata modification transactions using read committed snapshot isolation **414**. Differences between snapshot isolation **314** of FIG. **3** and read committed snapshot isolation **414** of FIG. **4** are described as follows.

[0058] With regard to snapshot isolation **314** and read committed snapshot isolation **414**, the term “snapshot” reflects the fact that all queries in a transaction see (are applied to) the same version, or “snapshot”, of the database based on the state of the database at the moment in time when the transaction begins. The term “isolation” refers to the fact that different transactions that see different versions of the database may nevertheless run concurrently because the snapshots effectively isolate such transactions from one another.

[0059] The difference between read committed snapshot isolation and snapshot isolation is that version visibility is scoped to the statement level in read committed snapshot isolation, whereas such visibility is scoped to the transaction level for snapshot isolation. This difference will be better appreciated through further consideration of FIGS. **3** and **4**. Timeline **300** of FIG. **3** is set forth as a table having rows and columns and wherein the time axis runs vertical up and down the table. That is, moving down from one row to the next is equivalent to moving forward in time (i.e., moving down the table means moving forward in time). Timeline **300** of FIG. **3** includes the following columns: metadata modification command **302**, lock type **304**, time **306**, data read command **308**, lock type **310** and comments **312**.

[0060] Metadata modification command **302** of FIG. **3** depicts a transaction that modifies metadata. In particular, the ALTER TABLE statement at time=**T6** causes a new column to be added to table named ‘Employees’ which causes the metadata object associated with table ‘Employees’ to be modified. The transaction for the ALTER TABLE command starts at time=**T4** and commits at time=**T9**.

[0061] Data read command **308** of FIG. **3** includes two transactions. The first transaction starts at time=**T1** and commits at time=**T12**, whereas the second transaction begins at time=**T13** and commits at time=**T15**. The first transaction includes two SELECT statements that begin executing at time=**T2** and time=**T10**, and which cause data and metadata to be read. Per snapshot isolation semantics, the SELECT statement that begins executing at time=**T2** does NOT see the new column ‘RemainingSickDays’ that is added by metadata modification command **302** because the change brought about by metadata modification command **302** does not exist until such time that its corresponding transaction commits at time=**T9**.

[0062] The second SELECT statement executes inside the first transaction at time=**T10**. Note, time=**T10** is AFTER metadata modification command **302** has committed so in theory, the modified metadata should be available for statements that execute after that point. With snapshot isolation, however, visibility of metadata modifications is scoped to the transaction level.

Accordingly, metadata versions that are visible cannot change during the course of a transaction, and the visible metadata version is fixed to the point in time when a transaction starts. With continued reference to FIG. 3, the second transaction that starts at time=T13 \*does\* see the modified metadata since that transaction started (at time=T13) well after the transaction of metadata modification command 302 committed (at time=T9).

[0063] In read committed snapshot isolation, on the other hand, visibility of metadata modifications is scoped to the statement level, and not the transaction level. This difference will now be described with reference to timeline 400 of FIG. 4 which includes the following columns: metadata modification command 402, lock type 404, time 406, data read command 408, lock type 410 and comments 412.

[0064] Metadata modification command 402 of FIG. 4 depicts a transaction that modifies metadata. In particular, the ALTER TABLE statement at time=T5 causes a new column to be added to table named 'Employees' which causes the metadata object associated with table 'Employees' to be modified. The transaction for the ALTER TABLE command starts at time=T3 and commits at time=T11.

[0065] Data read command 408 of FIG. 4 includes one transaction that executes three SELECT statements. The transaction starts at time=T1 and commits at time=T13. The first SELECT statement begins executing at time=T2 and completes executing at time=T9. The second select statement executes at time=T10, and the third SELECT statement executes at time=T12. Of course, the result of metadata modification command 402 cannot possibly be visible to any transaction or statement until such time as its transaction commits at time T=T11. Accordingly, the new column added by the ALTER TABLE statement at T=T5 is not visible to either of the first two SELECT statements that start at time=T2 and time=T10. However, the SELECT statement that executes at time=T12 does so \*AFTER\* metadata modification command transaction commits at time=T11. Because of this, and also because example 402 is using read committed snapshot isolation, the new column added by metadata modification command 402 \*is\* visible to the SELECT statement at time=T12. This is the case because in read committed snapshot isolation, visibility is determined for each statement in a transaction. In snapshot isolation, on the other hand, visibility is determined for an entire transaction and not for each statement within that transaction.

[0066] In an embodiment, the SELECT statement or data read command 408 that executes at time=T10 causes a query plan to be generated and then optimized by binding/algebraizer/query optimizer 112, and that query plan is subsequently executed by query execution engine 120. In an embodiment, the compiled query plan is cached and re-used for subsequent executions of the same statement. With continued reference to FIG. 4, although the SELECT statement that executes at time=T12 is the same statement as the one that executed at time=T10, the query plan compiled for the latter SELECT statement cannot be used because that plan refers to an old version of metadata that corresponds to the Employees table before the new column was added by the ALTER TABLE command at time=T5. Embodiments will detect that the cached query plan that was generated when the SELECT statement at time=T10 was executed is, in fact, stale since it refers to the wrong version of the metadata, and accordingly embodiments will re-compile the query plan to refer to the proper version(s) of metadata. The recompiled query plan may then be executed and doing so as depicted in data read command 408 at time=T12 will illustrate that the new column added by the ALTER TABLE of metadata modification command 402 is, in fact, visible.

[0067] Example logical timelines 300 and 400 of FIGS. 3 and 4, respectively, also include lock type columns 304, 310, 404 and 410. The lock type column reflects the lock type being requested and held by a statement inside the transaction. As described herein above, conventional lock types include a schema modification lock (SCH-M) and a schema stability lock (SCH-S). Embodiments of relational database system 100 also include and use these types of locks. This is true because some metadata modification commands (i.e., DDL commands) are not compatible with nor make use of versioned metadata, and too because some metadata operations should and/or must have

exclusive access to the metadata object.

[0068] For these reasons, embodiments implement two new lock types called access schema version and create schema version denoted by 'SCH-A' and 'SCH-C', respectively. These lock types are used for statements and DDL operations that are version aware and compatible (e.g., ALTER TABLE and/or SELECT, as described above). Due to the use of versioned metadata, there is no longer a need for full and exclusive access to a metadata object. A statement that requests and receives a SCH-A lock from, for example, lock manager **124** will automatically see the metadata versions visible to it when the statement begins to execute. Likewise, a statement like ALTER TABLE that requests and receives a SCH-C lock may change a metadata object without being blocked by a corresponding read when the result of the ALTER TABLE is a new version of the metadata object. Accordingly, there is no conflict between the locks, neither such lock blocks the other and the database is ACID compliant and maintained in a consistent state.

[0069] In embodiments, the schema modification lock (SCH-M) mentioned above is acquired when a DDL statement is executed and prevents access to the locked object data as the structure of the object is being changed. SQL Server allows a single schema modification lock (SCH-M) lock on any locked object. The schema stability lock (SCH-S) is acquired while a schema-dependent query is being compiled and executed and execution plan is generated. This lock does not block other transactions to access the object data.

[0070] With regard to SCH-A (access a schema version) and SCH-C (create a schema version), embodiments may support and/or enforce various types of constraints or functionality such as:

[0071] 1. Not all DDL operations may support metadata versioning [0072] 2. No concurrent DDL operations (i.e., only one DDL operation per object at a time) regardless of whether the operation is one that supports versioning. [0073] 2. Data manipulation language ("DML") operations such as, e.g., insert, update, delete or merge, are not allowed while schema changes are in progress i.e., no DML operation while there are pending (uncommitted) DDL changes. [0074] 3. Snapshot queries under snapshot isolation or read committed snapshot isolation are allowed while there are ongoing DDL operations. Other isolation levels are blocked. [0075] 4. Some operations cannot or will not support versioning and should block even snapshot queries. Such operations may include, for example, DBCC SHRINK. During DBCC SHRINK, operation affects physical structure of the index such as when moving index allocation map ("IAM") pages. As such, a SCH-M lock is acquired. Furthermore, a subset of DDL operations (mainly add column, remove column) may be supported and work incrementally. In an embodiment, any un-supported DDL operations (such as rename column etc.) may maintain the old blocking semantics.

[0076] How the new lock modes interact with existing schema locks: [0077] SCH-A: Indicates that the query requires schema stability using versioning. SCH-A is acquired by snapshot queries that currently take SCH-S and conflicts with SCH-M lock that will still be acquired by all DDL operations that don't support versioning. [0078] SCH-C: Indicates a versioned schema modification and is acquired by DDL operations that support schema versioning, for example ADD COLUMN. This is less restrictive than SCH-M in the sense that it is compatible with SCH-A, but still incompatible with SCH-S.

[0079] The lock modes regarding schema changes and stability may include the following, starting from the least to the most exclusive: [0080] SCH-A: Acquired by snapshot queries that used to take SCH-S i.e., all SI will take SCH-A lock. [0081] SCH-S: Acquired by all other operations that use it today, such as NOLOCK scans, etc.

[0082] Note that even though SCH-A may be considered the weakest lock (in new locking schema), SCH-S is still acquired for NOLOCK table scan to prevent schema changes while the scan is in progress. Data access locks, such as IS, IX, etc. are still used by other isolation levels, such as Read Committed, etc. SCH-C: Acquired by the DDL operations that support versioning but enables snapshot queries that now use SCH-A. SCH-M: Acquired by non-versioned DDL operations and blocks everything such as DBCC shrink, etc.

[0083] A compatibility matrix (C—Conflict, N—No conflict) is shown below as Table 1 for snapshot queries:

TABLE-US-00001	TABLE 1	Granted Mode	Requested Mode	SCH-S	SCH-M	SCH-A	SCH-C																	
S/U/X/IS/IU/IX	SCH-S	N	C	N	C	N	SCH-M	C	C	C	C	C	SCH-A	N	C	N	N	N	SCH-C	C	C	N	C	C
S/U/X/IS/IU/IX	N	C	N	C	various																			

[0084] Further changes may include: [0085] 1. In storage engine, acquire SCH-S when accessing HoBt (Heap or B-tree) metadata. With metadata versioning, switch snapshot access calls from SCH-S to SCH-A. [0086] 2. Similarly for system table level snapshot access, acquire SCH-A instead of SCH-S lock. [0087] 3. For query compilation, use RCSI semantics and use SCH-A lock instead of SCH-S. During execution time if the user transaction is not running at SI/RCSI level, then may recompile the query based on current object version.

[0088] FIGS. 5A, 5B, 5C and 6 which depict flowcharts 500, 510, 530 and 600, respectively, of example methods of operating a relational database system to modify a metadata object without being blocked by a concurrent read access of the metadata object, according to an embodiment. Flowcharts 500, 510, 530 and 600 illustrate methods that are analogous to the snapshot isolation and read committed snapshot isolation scenarios described herein above with regard to example logical timelines 300 and 400 of FIGS. 3 and 4, respectively. In an embodiment, flowcharts 500, 510, 530 and 600 may be performed by relational database system 100 of FIG. 1. Although described with reference to system 100 as shown in FIG. 1 and logical timeline 300 of FIG. 3, flowcharts 500, 510, 530 and 600 are not limited to that implementation. Other structural and operational embodiments will be apparent to persons skilled in the relevant art(s) based on the following discussion flowcharts 500, 510, 530 and 600 of FIGS. 5A, 5B, 5C and 6, respectively.

[0089] As shown in FIG. 5A, flowchart 500 begins with step 502. In step 502, a data read command is received that includes first and second queries, each query referencing a metadata object in metadata storage of the relational database system and configured to read user data objects from a user data storage of the relational database system. For example, client 104 of may send a read command to relational database system 100 as shown in FIG. 1. More specifically, client 104 may send data read command 308 or 408 as shown in FIGS. 3 and 4, respectively, which is received by parser 110 within execution environment 108. Data read command 308 and 408 each include first and second queries in the form of the SELECT statements directed at the Employees table that begin executing at time T2 and T10, respectively. Processing such queries within execution environment 108 requires that the metadata objects necessarily reference the metadata that corresponds to the Employees table. Flowchart 500 continues at step 504.

[0090] In step 504, a metadata modification command is received to be executed concurrently with the data read command, and wherein execution of the metadata modification command causes the metadata object stored in a metadata storage of the relational database system to be modified. For example, and with continued reference to FIGS. 1, 3 and 4, client 106 may send metadata modification command 302 or 402 as shown in FIG. 3 or 4, respectively, which is received by parser 110 within execution environment 108. Metadata modification command 302 and 402 include an ALTER TABLE command that begins executing at time T6 or T5, respectively, and which adds a RemainingSickDays column to the Employees table. Processing such the ALTER modifies the metadata that corresponds to the Employees table. Flowchart 500 continues at step 506.

[0091] In step 506, in response to receiving the data read command, the first query is executed in a first transaction. With continued reference to FIGS. 1, 3 and 4, data read command 308 or 408 executes each of the first and second queries within a transaction that begins at time T1 and commits at time T12 or T13, respectively. Flowchart 500 concludes at step 508.

[0092] In step 508, in response to receiving the metadata modification command, the metadata object is modified in a second transaction concurrently with the first transaction without being blocked by the first query executing in the first transaction. With continued reference to FIGS. 3

and **4**, the ALTER TABLE command that begins executing at time **T6** or **T5**, respectively, executes within a transaction that begins at time **T4** or **T5**, respectively, and commits at time **T9** or **T11**, respectively. Data read command **308** and **408** each obtains and uses a schema access lock SCH-A as denoted in column lock type **310** or **410** of FIGS. **3** and **4**, respectively. The ALTER TABLE command obtains and uses a schema change lock SCH-C as denoted in column lock type **304** and **404** of FIGS. **3** and **4**, respectively. As described herein above, the SCH-A lock does not block the SCH-C lock and the ALTER TABLE may execute concurrently with the queries being executed by data read command **308**.

[0093] In the foregoing discussion of steps **502-508** of flowchart **500**, it should be understood that at times, such steps may be performed in a different order or even contemporaneously with other steps. For example, the SELECT statement that begins executing at time **T2** could begin executing at a different time. Other operational embodiments will be apparent to persons skilled in the relevant art(s). Note also that the foregoing general description of the operation of relational database system **100** is provided for illustration only, and embodiments of relational database system **100** may comprise different hardware and/or software, and may operate in manners different than described above. Indeed, steps of flowchart **500** may be performed in various ways.

[0094] For example, FIG. **5B** depicts a flowchart **510** of an example method of operating a relational database system to modify a metadata object without being blocked by a concurrent read access of the metadata object, according to an embodiment, and wherein flowchart **510** comprises refinements or additions to the method steps of flowchart **500** as depicted in FIG. **5A**. Accordingly, flowchart **510** of FIG. **5B** will also be described with continued reference to relational database system **100** of FIG. **1** and FIGS. **3** and **4**. However, other structural and operational embodiments will be apparent to persons skilled in the relevant art(s) based on the following discussion regarding flowchart **510**.

[0095] As shown in FIG. **5B**, flowchart **510** begins with step **512**. In step **512**, a first transaction is started at a time **T1**. For example, and with reference to timeline **300** of FIG. **3**, data read command **308** starts a transaction at a time **T1** on the timeline denoted in time column **306**. Flowchart **510** continues at step **514**.

[0096] In step **514**, a metadata access lock is obtained from a lock manager of the relational database system. For example, The SELECT statement of data read command **308** that begins executing at time **T2** on the timeline denoted in time column **306** obtains a SCH-A lock from lock manager **124** within execution environment **108** of relational database system **100** as described further herein. Flowchart **510** continues at step **516**.

[0097] In step **516**, a first query plan is created for the first query, wherein the first query plan references the metadata object. With reference to FIG. **1**, when data read command **308** is processed within execution environment **108**, parser **110** in conjunction with binding/algebraizer/query optimizer **112** will cause a query plan to be created that references the proper metadata objects that correspond to the snapshot created when transaction begins. More specifically, a query plan for the SELECT statement of data read command **308** that begins executing at time **T2** on the timeline denoted in time column **306** will reference the metadata for the Employees table that exists in the snapshot created a time **T1** when the transaction begins. Flowchart **510** continues at step **518**.

[0098] In step **518**, execution of the first query plan is begun inside the first transaction at a time **T2** using the first metadata access lock. For example, the query plan for the SELECT statement of data read command **308** that begins executing at time **T2** on the timeline denoted in time column **306** is submitted to query execution engine **120** as depicted in FIG. **1**, and such query plan begins executing. Flowchart **510** concludes at step **516**.

[0099] In step **520**, execution of the first query plan is completed at a time **T3**. With continued reference to FIG. **3**, the SELECT statement of data read command **308** that begins executing at time **T2** on the timeline denoted in time column **306** completes execution at time **T9** as shown on

the timeline denoted in time column **306**. Thus, the time **T3** corresponds to the time **T9** as shown on the timeline denoted in time column **306**.

[0100] In step **522**, the first transaction is committed at a time **T4**. For example, the transaction that begun at time **T1** on the timeline denoted in time column **306** commits at time **T12** on that same timeline. Thus, the time **T4** in this example, is the time **T12** on the timeline denoted in time column **306**.

[0101] In the foregoing discussion of steps **512-522** of flowchart **510**, it should be understood that at times, such steps may be performed in a different order or even contemporaneously with other steps. For example, the **SELECT** statement that begins executing at time **T2** could begin executing at a different time. Other operational embodiments will be apparent to persons skilled in the relevant art(s). Note also that the foregoing general description of the operation of relational database system **100** is provided for illustration only, and embodiments of relational database system **100** may comprise different hardware and/or software, and may operate in manners different than described above.

[0102] FIG. **5C** depicts a flowchart **530** that is a refinement to the method depicted in FIGS. **5A** and **5B**, according to an embodiment. Accordingly, flowchart **530** of FIG. **5B** is also described with continued reference to relational database system **100** of FIG. **1** and FIGS. **3** and **4**. However, other structural and operational embodiments will be apparent to persons skilled in the relevant art(s) based on the following discussion regarding flowchart **530**.

[0103] As shown in FIG. **5C**, flowchart **530** begins with step **532**. In step **532**, a second transaction is started concurrently with the first transaction at a time **T5**. For example, and with reference to timeline **300** of FIG. **3**, metadata modification command **302** includes a transaction that begins at time **T4** as denoted on the timeline denoted in time column **306**. Thus, the time **T5** in flowchart **530** may correspond, for example, to time **T4** as denoted on the timeline denoted in time column **306**. Flowchart **530** continues with step **534**.

[0104] In step **534**, a metadata change lock is obtained from the lock manager. For example, The **ALTER TABLE** statement of metadata modification command **302** that begins executing at time **T6** on the timeline denoted in time column **306** obtains a SCH-C lock from lock manager **124** within execution environment **108** of relational database system **100** as described further herein above. Flowchart **530** continues at step **536**.

[0105] In step **536**, the metadata object is modified inside the second transaction using the metadata change lock wherein the modification is not blocked by the metadata access lock. With continued reference to FIG. **3**, the **ALTER TABLE** command that begins executing at time **T6** obtains a SCH-C lock whereas data read command **308** obtains and uses a schema access lock SCH-A as denoted in column lock type **310** of FIG. **3**. As described herein above, the SCH-A lock does not block the SCH-C lock and the **ALTER TABLE** may execute concurrently with the queries being executed by data read command **308**. Flowchart **530** concludes at step **538**.

[0106] In step **538**, the second transaction is committed at a time **T6**. With continued reference to FIG. **3**, the **ALTER TABLE** command that begins executing at time **T6** executes within a transaction that begins at time **T4** and commits at time **T9**. Thus, the time **T6** corresponds to the time **T9** as denoted on the timeline denoted in time column **306**.

[0107] In the foregoing discussion of steps **532-538** of flowchart **530**, it should be understood that at times, such steps may be performed in a different order or even contemporaneously with other steps. Other operational embodiments will be apparent to persons skilled in the relevant art(s). Note also that the foregoing general description of the operation of relational database system **100** is provided for illustration only, and embodiments of relational database system **100** may comprise different hardware and/or software, and may operate in manners different than described above. Indeed, steps of flowcharts **500**, **510** and **530** may be performed in various ways.

[0108] For example, flowchart **600** of FIG. **6** illustrates an example refinement to flowcharts **500**, **510** and/or **530** wherein embodiments operate in a read committed snapshot isolation mode. In such

a mode, metadata object version visibility and new snapshot creation is performed on for each command within a transaction rather than being scoped entirely to a transaction. Flowchart **600** begins with step **602**. In step **602**, the second query is executed inside the first transaction at a time prior to **T6** causing a second query plan to be compiled, wherein the second query plan references the first version of the metadata object. For example, and with reference to FIG. **4**, data read command **408** includes a first SELECT statement that begins executing at time **T2** as shown in time column **406**. That query will cause a query plan to be created for its execution. Likewise, the SELECT statement of data read command **408** that begins executing at time **T10** as shown in time column **406** will cause a different query plan to be created since the query is different. In each case, the transaction of metadata modification command **402** is NOT completed when such SELECT statements begin to execute. Accordingly, such queries will see a snapshot that references the version of the metadata of the Employees table that existed when such SELECT statements started. Flowchart **600** concludes at step **604**.

[0109] In step **604**, the second query is executed a second time inside the first transaction at a time subsequent to time **T6** causing the second query plan to be recompiled to generate a third query plan, wherein the third query plan references the second version of the metadata object. With continued reference to FIG. **4**, the SELECT statement that begins executing at time **T10** is executed again at time **T12**. However, because metadata modification command **402** transaction committed at time **T11**, the SELECT statement executing at time **T12** cannot use the same query plan because that query plan references an old version of the metadata corresponding to the Employees table. Instead, embodiments will create a new snapshot that references the second version of the metadata for the Employees table wherein the new RemainingSickDays column is in fact visible.

[0110] The scenario described in flowchart **600** is not applicable where embodiments are operating in the snapshot isolation mode such as depicted by example timeline **300** of FIG. **3**. In that instance, and as described above, the visibility of metadata modifications is always scoped to the transaction level rather than the statement level within a transaction. Accordingly, every statement that executes within, for example, the transaction that begins at time **T1** as shown in time column **306** of FIG. **3** will see the metadata version that existed at the time the transaction starts even though a concurrently executed metadata modification command may execute and commit before for data read command transaction commits.

[0111] In the foregoing discussion of the steps of flowcharts **500**, **510**, **530** and **600**, it should be understood that at times, such steps may be performed in a different order or even contemporaneously with other steps. Note also that the foregoing general description of the operation of relational database system **100** is provided for illustration only, and embodiments of relational database system **100** may comprise different hardware and/or software, and may operate in manners different than described above.

### III. Example Computer System Implementation

[0112] Each of relational database system **100**, client **104**, client **106**, binding/algebraizer/query optimizer **112**, parser **110**, lock manager **124**, metadata storage **114**, metadata cache **116**, data storage engine **118**, query execution engine **120**, and/or transaction manager **122**, and flowchart **500**, **510** and/or **530** may be implemented in hardware, or hardware combined with software and/or firmware. For example, relational database system **100**, client **104**, client **106**, binding/algebraizer/query optimizer **112**, parser **110**, lock manager **124**, metadata storage **114**, metadata cache **116**, data storage engine **118**, query execution engine **120**, and/or transaction manager **122**, and flowchart **500**, **510** and/or **530** may be implemented as computer program code/instructions configured to be executed in one or more processors and stored in a computer readable storage medium. Alternatively, relational database system **100**, client **104**, client **106**, binding/algebraizer/query optimizer **112**, parser **110**, lock manager **124**, metadata storage **114**, metadata cache **116**, data storage engine **118**, query execution engine **120**, and/or transaction manager **122**, and flowchart **500**, **510**, **530**, and/or **600** may be implemented as hardware

logic/electrical circuitry.

[0113] For instance, in an embodiment, one or more, in any combination, of relational database system **100**, client **104**, client **106**, binding/algebraizer/query optimizer **112**, parser **110**, lock manager **124**, metadata storage **114**, metadata cache **116**, data storage engine **118**, query execution engine **120**, and/or transaction manager **122**, and flowchart **500**, **510** and/or **530** may be implemented together in a SoC. The SoC may include an integrated circuit chip that includes one or more of a processor (e.g., a central processing unit (CPU), microcontroller, microprocessor, digital signal processor (DSP), one or more graphics processing units (GPUs), etc.), memory, one or more communication interfaces, and/or further circuits, and may optionally execute received program code and/or include embedded firmware to perform functions.

[0114] FIG. 7 depicts an exemplary implementation of a computing device **700** in which embodiments may be implemented. For example, relational database system **100**, client **104**, client **106**, binding/algebraizer/query optimizer **112**, parser **110**, lock manager **124**, metadata storage **114**, metadata cache **116**, data storage engine **118**, query execution engine **120**, and/or transaction manager **122** may be implemented in one or more computing devices similar to computing device **700** in stationary or mobile computer embodiments, including one or more features of computing device **700** and/or alternative features. The description of computing device **700** provided herein is provided for purposes of illustration, and is not intended to be limiting. Embodiments may be implemented in further types of computer systems, as would be known to persons skilled in the relevant art(s).

[0115] Computing device **702** can be any of a variety of types of computing devices. For example, computing device **702** may be a mobile computing device such as a handheld computer (e.g., a personal digital assistant (PDA)), a laptop computer, a tablet computer (such as an Apple iPad™), a hybrid device, a notebook computer (e.g., a Google Chromebook™ by Google LLC), a netbook, a mobile phone (e.g., a cell phone, a smart phone such as an Apple® iPhone® by Apple Inc., a phone implementing the Google® Android™ operating system, etc.), a wearable computing device (e.g., a head-mounted augmented reality and/or virtual reality device including smart glasses such as Google® Glass™, Oculus Rift® of Facebook Technologies, LLC, etc.), or other type of mobile computing device. Computing device **702** may alternatively be a stationary computing device such as a desktop computer, a personal computer (PC), a stationary server device, a minicomputer, a mainframe, a supercomputer, etc.

[0116] As shown in FIG. 7, computing device **702** includes a variety of hardware and software components, including a processor **710**, a storage **720**, one or more input devices **730**, one or more output devices **750**, one or more wireless modems **760**, one or more wired interfaces **780**, a power supply **782**, a location information (LI) receiver **784**, and an accelerometer **786**. Storage **720** includes memory **756**, which includes non-removable memory **722** and removable memory **724**, and a storage device **790**. Storage **720** also stores an operating system **712**, application programs **714**, and application data **716**. Wireless modem(s) **760** include a Wi-Fi modem **762**, a Bluetooth modem **764**, and a cellular modem **766**. Output device(s) **750** includes a speaker **752** and a display **754**. Input device(s) **730** includes a touch screen **732**, a microphone **734**, a camera **736**, a physical keyboard **738**, and a trackball **740**. Not all components of computing device **702** shown in FIG. 7 are present in all embodiments, additional components not shown may be present, and any combination of the components may be present in a particular embodiment. These components of computing device **702** are described as follows.

[0117] A single processor **710** (e.g., central processing unit (CPU), microcontroller, a microprocessor, signal processor, ASIC (application specific integrated circuit), and/or other physical hardware processor circuit) or multiple processors **710** may be present in computing device **702** for performing such tasks as program execution, signal coding, data processing, input/output processing, power control, and/or other functions. Processor **710** may be a single-core or multi-core processor, and each processor core may be single-threaded or multithreaded (to



provide multiple threads of execution concurrently). Processor **710** is configured to execute program code stored in a computer readable medium, such as program code of operating system **712** and application programs **714** stored in storage **720**. Operating system **712** controls the allocation and usage of the components of computing device **702** and provides support for one or more application programs **714** (also referred to as “applications” or “apps”). Application programs **714** may include common computing applications (e.g., e-mail applications, calendars, contact managers, web browsers, messaging applications), further computing applications (e.g., word processing applications, mapping applications, media player applications, productivity suite applications), one or more machine learning (ML) models, as well as applications related to the embodiments disclosed elsewhere herein.

[0118] Any component in computing device **702** can communicate with any other component according to function, although not all connections are shown for ease of illustration. For instance, as shown in FIG. 7, bus **706** is a multiple signal line communication medium (e.g., conductive traces in silicon, metal traces along a motherboard, wires, etc.) that may be present to communicatively couple processor **710** to various other components of computing device **702**, although in other embodiments, an alternative bus, further buses, and/or one or more individual signal lines may be present to communicatively couple components. Bus **706** represents one or more of any of several types of bus structures, including a memory bus or memory controller, a peripheral bus, an accelerated graphics port, and a processor or local bus using any of a variety of bus architectures.

[0119] Storage **720** is physical storage that includes one or both of memory **756** and storage device **790**, which store operating system **712**, application programs **714**, and application data **716** according to any distribution. Non-removable memory **722** includes one or more of RAM (random access memory), ROM (read only memory), flash memory, a solid-state drive (SSD), a hard disk drive (e.g., a disk drive for reading from and writing to a hard disk), and/or other physical memory device type. Non-removable memory **722** may include main memory and may be separate from or fabricated in a same integrated circuit as processor **710**. As shown in FIG. 7, non-removable memory **722** stores firmware **718**, which may be present to provide low-level control of hardware. Examples of firmware **718** include BIOS (Basic Input/Output System, such as on personal computers) and boot firmware (e.g., on smart phones). Removable memory **724** may be inserted into a receptacle of or otherwise coupled to computing device **702** and can be removed by a user from computing device **702**. Removable memory **724** can include any suitable removable memory device type, including an SD (Secure Digital) card, a Subscriber Identity Module (SIM) card, which is well known in GSM (Global System for Mobile Communications) communication systems, and/or other removable physical memory device type. One or more of storage device **790** may be present that are internal and/or external to a housing of computing device **702** and may or may not be removable. Examples of storage device **790** include a hard disk drive, a SSD, a thumb drive (e.g., a USB (Universal Serial Bus) flash drive), or other physical storage device.

[0120] One or more programs may be stored in storage **720**. Such programs include operating system **712**, one or more application programs **714**, and other program modules and program data. Examples of such application programs may include, for example, computer program logic (e.g., computer program code/instructions) for implementing one or more of relational database system **100**, client **104**, client **106**, binding/algebraizer/query optimizer **112**, parser **110**, lock manager **124**, metadata storage **114**, metadata cache **116**, data storage engine **118**, query execution engine **120**, and/or transaction manager **122**, and flowchart **500**, **510**, **530**, and/or **600** (including any suitable step of flowchart **500**, **510**, **530** and/or **600**), and/or further embodiments described herein.

[0121] Storage **720** also stores data used and/or generated by operating system **712** and application programs **714** as application data **716**. Examples of application data **716** include web pages, text, images, tables, sound files, video data, and other data, which may also be sent to and/or received from one or more network servers or other devices via one or more wired or wireless networks.

Storage **720** can be used to store further data including a subscriber identifier, such as an International Mobile Subscriber Identity (IMSI), and an equipment identifier, such as an International Mobile Equipment Identifier (IMEI). Such identifiers can be transmitted to a network server to identify users and equipment.

[0122] A user may enter commands and information into computing device **702** through one or more input devices **730** and may receive information from computing device **702** through one or more output devices **750**. Input device(s) **730** may include one or more of touch screen **732**, microphone **734**, camera **736**, physical keyboard **738** and/or trackball **740** and output device(s) **750** may include one or more of speaker **752** and display **754**. Each of input device(s) **730** and output device(s) **750** may be integral to computing device **702** (e.g., built into a housing of computing device **702**) or external to computing device **702** (e.g., communicatively coupled wired or wirelessly to computing device **702** via wired interface(s) **780** and/or wireless modem(s) **760**). Further input devices **730** (not shown) can include a Natural User Interface (NUI), a pointing device (computer mouse), a joystick, a video game controller, a scanner, a touch pad, a stylus pen, a voice recognition system to receive voice input, a gesture recognition system to receive gesture input, or the like. Other possible output devices (not shown) can include piezoelectric or other haptic output devices. Some devices can serve more than one input/output function. For instance, display **754** may display information, as well as operating as touch screen **732** by receiving user commands and/or other information (e.g., by touch, finger gestures, virtual keyboard, etc.) as a user interface. Any number of each type of input device(s) **730** and output device(s) **750** may be present, including multiple microphones **734**, multiple cameras **736**, multiple speakers **752**, and/or multiple displays **754**.

[0123] One or more wireless modems **760** can be coupled to antenna(s) (not shown) of computing device **702** and can support two-way communications between processor **710** and devices external to computing device **702** through network **704**, as would be understood to persons skilled in the relevant art(s). Wireless modem **760** is shown generically and can include a cellular modem **766** for communicating with one or more cellular networks, such as a GSM network for data and voice communications within a single cellular network, between cellular networks, or between the mobile device and a public switched telephone network (PSTN). Wireless modem **760** may also or alternatively include other radio-based modem types, such as a Bluetooth modem **764** (also referred to as a “Bluetooth device”) and/or Wi-Fi **762** modem (also referred to as an “wireless adaptor”). Wi-Fi modem **762** is configured to communicate with an access point or other remote Wi-Fi-capable device according to one or more of the wireless network protocols based on the IEEE (Institute of Electrical and Electronics Engineers) 802.11 family of standards, commonly used for local area networking of devices and Internet access. Bluetooth modem **764** is configured to communicate with another Bluetooth-capable device according to the Bluetooth short-range wireless technology standard(s) such as IEEE 802.15.1 and/or managed by the Bluetooth Special Interest Group (SIG).

[0124] Computing device **702** can further include power supply **782**, LI receiver **784**, accelerometer **786**, and/or one or more wired interfaces **780**. Example wired interfaces **780** include a USB port, IEEE **794** (FireWire) port, a RS-232 port, an HDMI (High-Definition Multimedia Interface) port (e.g., for connection to an external display), a DisplayPort port (e.g., for connection to an external display), an audio port, an Ethernet port, and/or an Apple® Lightning® port, the purposes and functions of each of which are well known to persons skilled in the relevant art(s). Wired interface(s) **780** of computing device **702** provide for wired connections between computing device **702** and network **704**, or between computing device **702** and one or more devices/peripherals when such devices/peripherals are external to computing device **702** (e.g., a pointing device, display **754**, speaker **752**, camera **736**, physical keyboard **738**, etc.). Power supply **782** is configured to supply power to each of the components of computing device **702** and may receive power from a battery internal to computing device **702**, and/or from a power cord plugged

into a power port of computing device **702** (e.g., a USB port, an A/C power port). LI receiver **784** may be used for location determination of computing device **702** and may include a satellite navigation receiver such as a Global Positioning System (GPS) receiver or may include other type of location determiner configured to determine location of computing device **702** based on received information (e.g., using cell tower triangulation, etc.). Accelerometer **786** may be present to determine an orientation of computing device **702**.

[0125] Note that the illustrated components of computing device **702** are not required or all-inclusive, and fewer or greater numbers of components may be present as would be recognized by one skilled in the art. For example, computing device **702** may also include one or more of a gyroscope, barometer, proximity sensor, ambient light sensor, digital compass, etc. Processor **710** and memory **756** may be co-located in a same semiconductor device package, such as being included together in an integrated circuit chip, FPGA, or system-on-chip (SOC), optionally along with further components of computing device **702**.

[0126] In embodiments, computing device **702** is configured to implement any of the above-described features of flowcharts herein. Computer program logic for performing any of the operations, steps, and/or functions described herein may be stored in storage **720** and executed by processor **710**.

[0127] In some embodiments, server infrastructure **770** may be present in computing environment **700** and may be communicatively coupled with computing device **702** via network **704**. Server infrastructure **770**, when present, may be a network-accessible server set (e.g., a cloud-based environment or platform). As shown in FIG. 7, server infrastructure **770** includes clusters **772**. Each of clusters **772** may comprise a group of one or more compute nodes and/or a group of one or more storage nodes. For example, as shown in FIG. 7, cluster **772** includes nodes **774**. Each of nodes **774** are accessible via network **704** (e.g., in a “cloud-based” embodiment) to build, deploy, and manage applications and services. Any of nodes **774** may be a storage node that comprises a plurality of physical storage disks, SSDs, and/or other physical storage devices that are accessible via network **704** and are configured to store data associated with the applications and services managed by nodes **774**. For example, as shown in FIG. 7, nodes **774** may store application data **778**.

[0128] Each of nodes **774** may, as a compute node, comprise one or more server computers, server systems, and/or computing devices. For instance, a node **774** may include one or more of the components of computing device **702** disclosed herein. Each of nodes **774** may be configured to execute one or more software applications (or “applications”) and/or services and/or manage hardware resources (e.g., processors, memory, etc.), which may be utilized by users (e.g., customers) of the network-accessible server set. For example, as shown in FIG. 7, nodes **774** may operate application programs **776**. In an implementation, a node of nodes **774** may operate or comprise one or more virtual machines, with each virtual machine emulating a system architecture (e.g., an operating system), in an isolated manner, upon which applications such as application programs **776** may be executed.

[0129] In an embodiment, one or more of clusters **772** may be co-located (e.g., housed in one or more nearby buildings with associated components such as backup power supplies, redundant data communications, environmental controls, etc.) to form a data center, or may be arranged in other manners. Accordingly, in an embodiment, one or more of clusters **772** may be a data center in a distributed collection of data centers. In embodiments, exemplary computing environment **700** comprises part of a cloud-based platform such as Amazon Web Services® of Amazon Web Services, Inc., or Google Cloud Platform™ of Google LLC, although these are only examples and are not intended to be limiting.

[0130] In an embodiment, computing device **702** may access application programs **776** for execution in any manner, such as by a client application and/or a browser at computing device **702**. Example browsers include Microsoft Edge® by Microsoft Corp. of Redmond, Washington, Mozilla

Firefox®, by Mozilla Corp. of Mountain View, California, Safari®, by Apple Inc. of Cupertino, California, and Google® Chrome by Google LLC of Mountain View, California.

[0131] For purposes of network (e.g., cloud) backup and data security, computing device **702** may additionally and/or alternatively synchronize copies of application programs **714** and/or application data **716** to be stored at network-based server infrastructure **770** as application programs **776** and/or application data **778**. For instance, operating system **712** and/or application programs **714** may include a file hosting service client, such as Microsoft® OneDrive® by Microsoft Corporation, Amazon Simple Storage Service (Amazon S3)® by Amazon Web Services, Inc., Dropbox® by Dropbox, Inc., Google Drive™ by Google LLC, etc., configured to synchronize applications and/or data stored in storage **720** at network-based server infrastructure **770**.

[0132] In some embodiments, on-premises servers **792** may be present in computing environment **700** and may be communicatively coupled with computing device **702** via network **704**. On-premises servers **792**, when present, are hosted within an organization's infrastructure and, in many cases, physically onsite of a facility of that organization. On-premises servers **792** are controlled, administered, and maintained by IT (Information Technology) personnel of the organization or an IT partner to the organization. Application data **798** may be shared by on-premises servers **792** between computing devices of the organization, including computing device **702** (when part of an organization) through a local network of the organization, and/or through further networks accessible to the organization (including the Internet). Furthermore, on-premises servers **792** may serve applications such as application programs **796** to the computing devices of the organization, including computing device **702**. Accordingly, on-premises servers **792** may include storage **794** (which includes one or more physical storage devices such as storage disks and/or SSDs) for storage of application programs **796** and application data **798** and may include one or more processors for execution of application programs **796**. Still further, computing device **702** may be configured to synchronize copies of application programs **714** and/or application data **716** for backup storage at on-premises servers **792** as application programs **796** and/or application data **798**.

[0133] Embodiments described herein may be implemented in one or more of computing device **702**, network-based server infrastructure **770**, and on-premises servers **792**. For example, in some embodiments, computing device **702** may be used to implement systems, clients, or devices, or components/subcomponents thereof, disclosed elsewhere herein. In other embodiments, a combination of computing device **702**, network-based server infrastructure **770**, and/or on-premises servers **792** may be used to implement the systems, clients, or devices, or components/subcomponents thereof, disclosed elsewhere herein.

[0134] As used herein, the terms “computer program medium,” “computer-readable medium,” and “computer-readable storage medium,” etc., are used to refer to physical hardware media. Examples of such physical hardware media include any hard disk, optical disk, SSD, other physical hardware media such as RAMs, ROMs, flash memory, digital video disks, zip disks, MEMs (microelectronic machine) memory, nanotechnology-based storage devices, and further types of physical/tangible hardware storage media of storage **720**. Such computer-readable media and/or storage media are distinguished from and non-overlapping with communication media and propagating signals (do not include communication media and propagating signals). Communication media embodies computer-readable instructions, data structures, program modules or other data in a modulated data signal such as a carrier wave. The term “modulated data signal” means a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, communication media includes wireless media such as acoustic, RF, infrared, and other wireless media, as well as wired media. Embodiments are also directed to such communication media that are separate and non-overlapping with embodiments directed to computer-readable storage media.

[0135] As noted above, computer programs and modules (including application programs **714**) may

be stored in storage **720**. Such computer programs may also be received via wired interface(s) **780** and/or wireless modem(s) **760** over network **704**. Such computer programs, when executed or loaded by an application, enable computing device **702** to implement features of embodiments discussed herein. Accordingly, such computer programs represent controllers of the computing device **702**.

[0136] Embodiments are also directed to computer program products comprising computer code or instructions stored on any computer-readable medium or computer-readable storage medium. Such computer program products include the physical storage of storage **720** as well as further physical storage types.

#### IV. Additional Example Embodiments

[0137] As described, systems and devices embodying the techniques herein may be configured and enabled in various ways to perform their respective functions for enabling metadata and data modification concurrency a relational database system. In embodiments, one or more of the steps or operations of any flowchart and/or flow diagram described herein may not be performed.

Moreover, steps or operations in addition to or in lieu of those in any flowchart and/or flow diagram described herein may be performed. Further, in examples, one or more operations of any flowchart and/or flow diagram described herein may be performed out of order, in an alternate sequence, or partially (or completely) concurrently with each other or with other operations.

[0138] As described herein, systems, devices, components, etc., of the embodiments that are configured to perform functions and/or operations are also contemplated as performing such functions and/or operations.

[0139] The additional examples and embodiments described in this Section may be applicable to examples disclosed in any other Section or subsection of this disclosure.

[0140] A relational database system is provided herein. The relational database system comprising: a metadata version aware execution engine configured to: receive a data read command comprising first and second queries, each query referencing a metadata object in metadata storage of the relational database system and configured to read user data objects from a user data storage of the relational database system; receive a metadata modification command to be executed concurrently with the data read command, and wherein execution of the metadata modification command causes the metadata object stored in a metadata storage of the relational database system to be modified; in response to receiving the data read command: start a first transaction at a time **T1**; obtain a metadata access lock from a lock manager of the relational database system; create a first query plan for the first query, wherein the first query plan references the metadata object; begin execution of the first query plan inside the first transaction at a time **T2** using the metadata access lock; complete execution of the first query plan at a time **T3**; and commit the first transaction at a time **T4**; and in response to receiving the metadata modification command: start a second transaction concurrently with the first transaction at a time **T5**; obtain a metadata change lock from the lock manager; modify the metadata object inside the second transaction using the metadata change lock, wherein the modification is not blocked by the metadata access lock; and commit the second transaction at a time **T6**.

[0141] In another embodiment of the relational database system, time **T2** is subsequent to time **T1**, time **T5** is subsequent to time **T2**, time **T3** is subsequent to time **T5** and time **T6** is subsequent to time **T3**.

[0142] In another embodiment of the relational database system, prior to executing the metadata modification command, the metadata object comprises a first version of the metadata object.

[0143] In another embodiment of the relational database system, executing the metadata modification command causes the first version of the metadata object to be modified into a second version of the metadata object.

[0144] In another embodiment of the relational database system, wherein the execution engine is configured to operate in a read committed snapshot isolation mode to: execute the second query

inside the first transaction at a time prior to T6 causing a second query plan to be compiled, wherein the second query plan references the first version of the metadata object; and execute the second query a second time inside the first transaction at a time subsequent to time T6 causing the second query plan to be recompiled to generate a third query plan, wherein the third query plan references the second version of the metadata object.

[0145] In another embodiment of the relational database system, wherein the execution engine is configured to operate in a regular snapshot isolation mode to: execute the second query inside the first transaction at a time either prior to or subsequent to time T6 causes a second query plan to be compiled wherein the second query plan references the first version of the metadata object; and execute the second query a second time inside the first transaction at a time subsequent to time T6 reusing the second query plan, wherein the query plan references the first version of the metadata object.

[0146] In another embodiment of the relational database system, the relational database system further comprises a transaction manager configured to determine which of the first or second version of the metadata object is visible to a transaction at a predetermined time.

[0147] In another embodiment of the relational database system, the transaction manager is further configured to determine which of the first or second version of the metadata object is visible to a transaction at a predetermined time based at least in part on whether the execution engine is configured to operate in a first snapshot isolation mode, a second snapshot isolation mode or a non-snapshot isolation mode.

[0148] A method in a relational database system is provided herein. The method comprising: receiving a data read command comprising first and second queries, each query referencing a metadata object in a metadata storage of the relational database system and configured to read user data objects from a user data storage of the relational database system; receiving a metadata modification command to be executed concurrently with the data read command, and wherein execution of the metadata modification command causes the metadata object to be modified; in response to receiving the data read command: starting a first transaction at a time T1; obtaining a metadata access lock from a lock manager of the relational database system; creating a first query plan for the first query, wherein the first query plan references the metadata object; beginning execution of the first query plan inside the first transaction at a time T2 using the metadata access lock; completing execution of the first query plan at a time T3; and committing the first transaction at a time T4; and in response to receiving the metadata modification command: starting a second transaction concurrently with the first transaction at a time T5; obtaining a metadata change lock from the lock manager; modifying the metadata object inside the second transaction using the metadata change lock, wherein the modification is not blocked by the metadata access lock; and committing the second transaction at a time T6.

[0149] In another embodiment of the foregoing method, time T2 is subsequent to time T1, time T5 is subsequent to time T2, time T3 is subsequent to time T5 and time T6 is subsequent to time T3.

[0150] In another embodiment of the foregoing method, prior to executing the metadata modification command, the metadata object comprises a first version of the metadata object.

[0151] In another embodiment of the foregoing method, executing the metadata modification command causes the first version of the metadata object to be modified into a second version of the metadata object.

[0152] In another embodiment of the foregoing method, the method further comprises operating the relational database system in a first snapshot isolation mode and: executing the second query inside the first transaction at a time prior to T6 causing a second query plan to be compiled, wherein the second query plan references the first version of the metadata object; and executing the second query a second time inside the first transaction at a time subsequent to time T6 causing the second query plan to be recompiled to generate a third query plan, wherein the third query plan references the second version of the metadata object.

[0153] In another embodiment of the foregoing method, the method further comprises operating the relational database system in a second snapshot isolation mode and executing the second query inside the first transaction at a time either prior to or subsequent to time T6 causing a second query plan to be compiled wherein the second query plan references the first version of the metadata object.

[0154] A computer program product comprising a computer-readable memory device having computer program logic recorded thereon that when executed by at least one processor of a computing device causes the at least one processor to perform operations is provided herein. The operations comprising: receiving a data read command comprising first and second queries, each query referencing a metadata object in a metadata storage of the relational database system and configured to read user data objects from a user data storage of the relational database system; receiving a metadata modification command to be executed concurrently with the data read command, and wherein execution of the metadata modification command causes the metadata object to be modified; in response to receiving the data read command: starting a first transaction at a time T1; obtaining a metadata access lock from a lock manager of the relational database system; creating a first query plan for the first query, wherein the first query plan references the metadata object; begin executing the first query plan inside the first transaction at a time T2 using the metadata access lock; completing execution of the first query plan at a time T3; and committing the first transaction at a time T4; and in response to receiving the metadata modification command: starting a second transaction concurrently with the first transaction at a time T5; obtaining a metadata change lock; modifying the metadata object inside the second transaction using the metadata change lock, wherein the modification is not blocked by the metadata access lock; and committing the second transaction at a time T6.

[0155] In another embodiment of the foregoing computer program product, time T2 is subsequent to time T1, time T5 is subsequent to time T2, time T3 is subsequent to time T5 and time T6 is subsequent to time T3.

[0156] In another embodiment of the foregoing computer program product, prior to executing the metadata modification command, the metadata object comprises a first version of the metadata object.

[0157] In another embodiment of the foregoing computer program product, executing the metadata modification command causes the first version of the metadata object to be modified into a second version of the metadata object.

[0158] In another embodiment of the foregoing computer program product, the operations further comprise operating the relational database system in a first snapshot isolation mode and: executing the second query inside the first transaction at a time prior to T6 causing a second query plan to be compiled, wherein the second query plan references the first version of the metadata object; and executing the second query a second time inside the first transaction at a time subsequent to time T6 causing the second query plan to be recompiled to generate a third query plan, wherein the third query plan references the second version of the metadata object.

[0159] In another embodiment of the foregoing computer program product, the operations further comprise operating the relational database system in the second snapshot isolation mode and executing the second query inside the first transaction at a time either prior to or subsequent to time T6 causing a second query plan to be compiled wherein the second query plan references the first version of the metadata object.

## V. Conclusion

[0160] References in the specification to “one embodiment,” “an embodiment,” “an example embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in

connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to affect such feature, structure, or characteristic in connection with other embodiments whether or not explicitly described.

[0161] In the discussion, unless otherwise stated, adjectives modifying a condition or relationship characteristic of a feature or features of an implementation of the disclosure, should be understood to mean that the condition or characteristic is defined to within tolerances that are acceptable for operation of the implementation for an application for which it is intended. Furthermore, if the performance of an operation is described herein as being “in response to” one or more factors, it is to be understood that the one or more factors may be regarded as a sole contributing factor for causing the operation to occur or a contributing factor along with one or more additional factors for causing the operation to occur, and that the operation may occur at any time upon or after establishment of the one or more factors. Still further, where “based on” is used to indicate an effect being a result of an indicated cause, it is to be understood that the effect is not required to only result from the indicated cause, but that any number of possible additional causes may also contribute to the effect. Thus, as used herein, the term “based on” should be understood to be equivalent to the term “based at least on.”

[0162] Numerous example embodiments have been described above. Any section/subsection headings provided herein are not intended to be limiting. Embodiments are described throughout this document, and any type of embodiment may be included under any section/subsection. Furthermore, embodiments disclosed in any section/subsection may be combined with any other embodiments described in the same section/subsection and/or a different section/subsection in any manner.

[0163] In some example embodiments, one or more of the operations of the flowcharts described herein may not be performed. Moreover, operations in addition to or in lieu of the operations of the flowcharts described herein may be performed. Further, in some example embodiments, one or more of the operations of the flowcharts described herein may be performed out of order, in an alternate sequence, or partially (or completely) concurrently with each other or with other operations.

[0164] While various embodiments have been described above, it should be understood that they have been presented by way of example only, and not limitation. It will be apparent to persons skilled in the relevant art that various changes in form and detail can be made therein without departing from the spirit and scope of the embodiments. Thus, the breadth and scope of the embodiments should not be limited by any of the above-described example embodiments, but should be defined only in accordance with the following claims and their equivalents.

## Claims

1. A relational database system, comprising: a processor; and a memory that stores program code structured to cause the processor to: begin a first transaction for a first query directed to first data and referencing a metadata object that is metadata to the first data, a first version of the metadata object stored in storage, the begin of the first transaction including obtaining a metadata access lock on the metadata object, and beginning execution of a first query plan for the first query; perform a second transaction subsequent to the begin of the first transaction at least by obtaining a metadata change lock on the metadata object, modifying the metadata object to create a second version of the metadata object using the metadata change lock and as allowed by the metadata access lock, storing the second version of the metadata object in the storage, and committing the second transaction; and complete the first transaction subsequent to the committing of the second transaction at least by completing execution of the first query plan, and committing the first transaction.
2. The relational database system of claim 1, wherein the program code is further structured to cause the processor to: execute a second query plan for a second query inside the first transaction at



a time prior to said committing the second transaction, wherein the second query plan references the first version of the metadata object; and recompile and execute the second query plan inside the first transaction at a time subsequent to said committing the second transaction, wherein the recompiled and executed query plan references the second version of the metadata object.

**3.** The relational database system of claim 1, wherein the program code is further structured to cause the processor to: execute a second query plan for a second query inside the first transaction a first time either prior to or subsequent to said committing the second transaction, wherein the second query plan references the first version of the metadata object; and execute the second query plan a second time inside the first transaction at a time subsequent to said committing the second transaction.

**4.** The relational database system of claim 1, wherein the program code comprises: a transaction manager configured to determine which of the first version or the second version of the metadata object is visible to a third transaction.

**5.** The relational database system of claim 4, wherein the transaction manager is further configured to determine which of the first version or the second version of the metadata object is visible to the third transaction based at least in part on whether the program code is configured to operate in a read committed snapshot isolation mode, a non-read committed snapshot isolation mode, or a non-snapshot isolation mode.

**6.** The relational database system of claim 4, wherein the first version of the metadata object is stored with a first transaction identifier of the first transaction and a first commit timestamp indicative of a first time at which the first version came into existence; The second version of the metadata object is stored with a second transaction identifier of the second transaction and a second commit timestamp indicative of a second time at which the second version came into existence; and wherein the transaction manager uses at least the first and second commit timestamps to determine which of the first version or the second version of the metadata object is visible to the third transaction.

**7.** The relational database system of claim 1, wherein a first statement of a third transaction that requests and receives the metadata access lock is enabled to automatically see metadata versions of a second metadata object that the first statement accesses; and Wherein a second statement of the third transaction that requests and receives the metadata change lock is enabled to change the second metadata object that the second statement accesses without being blocked by the metadata access lock of the first statement.

**8.** A method in a relational database system, comprising: beginning a first transaction for a first query directed to first data and referencing a metadata object that is metadata to the first data, a first version of the metadata object stored in storage, said beginning including obtaining a metadata access lock on the metadata object, and beginning execution of a first query plan for the first query; performing a second transaction subsequent to said beginning the first transaction at least by obtaining a metadata change lock on the metadata object, modifying the metadata object to create a second version of the metadata object using the metadata change lock and as allowed by the metadata access lock, storing the second version of the metadata object in the storage, and committing the second transaction; and completing the first transaction subsequent to the committing of the second transaction at least by completing execution of the first query plan, and committing the first transaction.

**9.** The method of claim 8, further comprising: executing a second query plan for a second query inside the first transaction at a time prior to said committing the second transaction, wherein the second query plan references the first version of the metadata object; and recompiling and executing the second query plan inside the first transaction at a time subsequent to said committing the second transaction, wherein the recompiled and executed query plan references the second version of the metadata object.

**10.** The method of claim 8, further comprising: executing a second query plan for a second query

inside the first transaction a first time either prior to or subsequent to said committing the second transaction, wherein the second query plan references the first version of the metadata object; and executing the second query plan a second time inside the first transaction at a time subsequent to said committing the second transaction.

**11.** The method of claim 8, further comprising: determining which of the first version or the second version of the metadata object is visible to a third transaction.

**12.** The method of claim 11, wherein said determining comprises: determining which of the first version or the second version of the metadata object is visible to the third transaction based at least in part on whether the program code is configured to operate in a read committed snapshot isolation mode, a non-read committed snapshot isolation mode, or a non-snapshot isolation mode.

**13.** The method of claim 11, wherein the first version of the metadata object is stored with a first transaction identifier of the first transaction and a first commit timestamp indicative of a first time at which the first version came into existence; the second version of the metadata object is stored with a second transaction identifier of the second transaction and a second commit timestamp indicative of a second time at which the second version came into existence; and the method further comprising: using at least the first and second commit timestamps to determine which of the first version or the second version of the metadata object is visible to the third transaction.

**14.** The method of claim 8, wherein a first statement of a third transaction that requests and receives the metadata access lock is enabled to automatically see metadata versions of a second metadata object that the first statement accesses; and wherein a second statement of the third transaction that requests and receives the metadata change lock is enabled to change the second metadata object that the second statement accesses without being blocked by the metadata access lock of the first statement.

**15.** A computer program product comprising a computer-readable memory device having computer program logic recorded thereon that when executed by at least one processor of a computing device causes the at least one processor to perform operations, the operations comprising: beginning a first transaction for a first query directed to first data and referencing a metadata object that is metadata to the first data, a first version of the metadata object stored in storage, said beginning including obtaining a metadata access lock on the metadata object, and beginning execution of a first query plan for the first query; performing a second transaction subsequent to said beginning the first transaction at least by obtaining a metadata change lock on the metadata object, modifying the metadata object to create a second version of the metadata object using the metadata change lock and as allowed by the metadata access lock, storing the second version of the metadata object in the storage, and committing the second transaction; and completing the first transaction subsequent to the committing of the second transaction at least by completing execution of the first query plan, and committing the first transaction.

**16.** The computer program product of claim 15, wherein the operations further comprise: executing a second query plan for a second query inside the first transaction at a time prior to said committing the second transaction, wherein the second query plan references the first version of the metadata object; and recompiling and executing the second query plan inside the first transaction at a time subsequent to said committing the second transaction, wherein the recompiled and executed query plan references the second version of the metadata object.

**17.** The computer program product of claim 15, wherein the operations further comprise: executing a second query plan for a second query inside the first transaction a first time either prior to or subsequent to said committing the second transaction, wherein the second query plan references the first version of the metadata object; and executing the second query plan a second time inside the first transaction at a time subsequent to said committing the second transaction.

**18.** The computer program product of claim 15, wherein the operations further comprise: determining which of the first version or the second version of the metadata object is visible to a third transaction.

**19.** The computer program product of claim 18, wherein said determining comprises: determining which of the first version or the second version of the metadata object is visible to the third transaction based at least in part on whether the program code is configured to operate in a read committed snapshot isolation mode, a non-read committed snapshot isolation mode, or a non-snapshot isolation mode.

**20.** The computer program product of claim 18, wherein the first version of the metadata object is stored with a first transaction identifier of the first transaction and a first commit timestamp indicative of a first time at which the first version came into existence; the second version of the metadata object is stored with a second transaction identifier of the second transaction and a second commit timestamp indicative of a second time at which the second version came into existence; and the operations further comprising: using at least the first and second commit timestamps to determine which of the first version or the second version of the metadata object is visible to the third transaction.

---