

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258840

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

ZHANG; Yang et al.

Data discovery and classification using sampling-based scanning techniques

Abstract

The technology disclosed herein relates to data discovery in computing environments, that provide access to data. In one example, a computer-implemented method includes identifying a first sampling criterion used for classification of a data store with respect to a target data type during a previous scan of the data store. The data store stores a set of data objects. The method includes selecting a second sampling criterion, from a plurality of sampling criteria, based on the first sampling criterion, and deploying one or more scanners configured to select a subset of data objects, from the set of data objects stored in the data store, based on the second sampling criterion. The subset comprises some, but not all, of the set of data objects. The method includes generating a classification result based on a number of instances of the target data type in the subset of data objects.

Inventors: ZHANG; Yang (Fremont, CA), Agrawal; Ajay (Bangalore, IN), ITHAL; Ravishankar Ganesh (Los Altos, CA)

Applicant: Proofpoint, Inc. (Sunnyvale, CA)

Family ID: 94372012

Assignee: Proofpoint, Inc. (Sunnyvale, CA)

Appl. No.: 19/193340

Filed: April 29, 2025

Foreign Application Priority Data

IN 202311050874

Jul. 28, 2023

Related U.S. Application Data

Publication Classification

Int. Cl.: G06F16/28 (20190101); H04L67/10 (20220101)

U.S. Cl.:

CPC G06F16/285 (20190101); H04L67/10 (20130101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] The present application is a continuation of and claims priority to U.S. patent application Ser. No. 18/754,643, filed Jun. 26, 2024, which is based on and claims the benefit of Indian Application No. 202311050874, filed Jul. 28, 2023; the contents of these applications are hereby incorporated by reference in their entirety.

FIELD OF THE TECHNOLOGY DISCLOSED

[0002] The technology disclosed generally relates to data discovery in computing environments, such as cloud environments, that provide access to storage resources that store data for users or other endpoints or systems. More specifically, but not by limitation, the present disclosure relates to improved systems and methods of data discovery in cloud or other computing environments for cloud security posture management (CSPM), cloud infrastructure entitlement management (CIEM), cloud-native application protection platform (CNAPP), cloud-native configuration management database (CMDB), and/or data security posture management (DSPM).

BACKGROUND

[0003] The subject matter discussed in this section should not be assumed to be prior art merely as a result of its mention in this section. Similarly, a problem mentioned in this section or associated with the subject matter provided as background should not be assumed to have been previously recognized in the prior art. The subject matter in this section merely represents different approaches, which in and of themselves can also correspond to implementations of the claimed technology.

[0004] There are many types of computing environments that provide data storage resources for organizations or other end users. Cloud computing, for example, provides on-demand availability of computer resources, such as data storage and compute resources, often without direct active management by users. Thus, a cloud environment can provide computation, software, data access, and storage services that do not require end-user knowledge of the physical location or configuration of the system that delivers the services. In various examples, remote servers can deliver the services over a wide area network, such as the Internet, using appropriate protocols, and those services can be accessed through a web browser or any other computing component.

[0005] Examples of cloud storage services include Amazon Web Services™ (AWS), Google Cloud Platform™ (GCP), and Microsoft Azure™, to name a few. Such cloud storage services provide on-demand network access to a shared pool of configurable resources. These resources can include networks, servers, storage, applications, services, etc. The end-users of such cloud services often include organizations that have a need to store sensitive and/or confidential data, such as personal information, financial information, medical information. Such information can be accessed by any of a number of users through permissions and access control data assigned or otherwise defined through administrator accounts.

[0006] The discussion above is merely provided for general background information and is not intended to be used as an aid in determining the scope of the claimed subject matter.

SUMMARY

[0007] The technology disclosed relates to data discovery in computing environments, such as cloud environments, that provide access to storage resources that store data for users or other endpoints or systems. The data discovery can identify sensitive and/or private data (or information), which can be used for data posture analysis and management. In one example, a computer-implemented method for classifying a data store in a computing environment includes identifying a first sampling criterion used for classification of the data store with respect to a target data type during a previous scan of the data store. The data store stores a set of data objects. The method also includes selecting a second sampling criterion, from a plurality of sampling criteria, based on the first sampling criterion used during the previous scan, and deploying one or more scanners configured to select a subset of data objects, from the set of data objects stored in the data store, based on the second sampling criterion. The subset comprises some, but not all, of the set of data objects. The method also includes generating a classification result based on a number of instances of the target data type in the subset of data objects and performing a computing action based on the classification result.

[0008] This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used as an aid in determining the scope of the claimed subject matter. The claimed subject matter is not limited to implementations that solve any or all disadvantages noted in the background.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0009] In the drawings, like reference characters generally refer to like parts throughout the different views. Also, the drawings are not necessarily to scale, with an emphasis instead generally being placed upon illustrating the principles of the technology disclosed. In the following description, various implementations of the technology disclosed are described with reference to the following drawings, in which:

[0010] FIG. 1 is a block diagram illustrating one example of a cloud architecture.

[0011] FIG. 2 is a block diagram illustrating one example of a cloud service.

[0012] FIG. 3 is a block diagram illustrating one example of a cloud security posture analysis system.

[0013] FIG. 4 is a block diagram illustrating one example of a deployed scanner.

[0014] FIG. 5 is a flow diagram showing an example operation of on-boarding a cloud account and deploying one or more scanners.

[0015] FIG. 6 illustrates one example of a user interface display having a dashboard representing on-boarded cloud service accounts.

[0016] FIG. 7 is a flow diagram illustrating one example of cloud infrastructure scanning performed by a cloud scanner deployed in a cloud service.

[0017] FIGS. 8-1 and 8-2 (collectively referred to as FIG. 8) provide a flow diagram illustrating one example of infrastructure analysis and query execution.

[0018] FIG. 9 illustrates an example user interface display.

[0019] FIGS. 10-1 and 10-2 (collectively referred to as FIG. 10) provide a flow diagram of an example operation of infrastructure posture analysis using incremental change detection.

[0020] FIG. 11 is a flow diagram illustrating one example of invoking an incremental change detector.

[0021] FIG. 12 is a flow diagram illustrating one example of filtering and applying risk signatures.
[0022] FIG. 13 illustrates one example of a user interface display that renders a set of tagged risk signatures.

[0023] FIG. 14 illustrates one example of a user interface display that renders instances of matched risk signatures.

[0024] FIG. 15 is a schematic diagram illustrating one example of a cloud data scanner deployment and control component.

[0025] FIG. 16 is a flow diagram illustrating one example of configuring and deploying a data scanner on a data store using sampling criterion.

[0026] FIG. 17 is a simplified block diagram of one example of a client device.

[0027] FIG. 18 illustrates an example of a handheld or mobile device.

[0028] FIG. 19 shows an example computer system.

DETAILED DESCRIPTION

[0029] The following discussion is presented to enable any person skilled in the art to make and use the technology disclosed, and is provided in the context of a particular application and its requirements. Various modifications to the disclosed implementations will be readily apparent to those skilled in the art, and the general principles defined herein may be applied to other implementations and applications without departing from the spirit and scope of the technology disclosed. Thus, the technology disclosed is not intended to be limited to the implementations shown, but is to be accorded the widest scope consistent with the principles and features disclosed herein.

[0030] As noted above, computing environments, such as cloud environments and/or on-premise environments, are used by organizations or other end-users to store a wide variety of different types of information in many contexts and for many uses. This data can often include sensitive and/or confidential information, and can be the target for malicious activity such as acts of fraud, privacy breaches, data theft, etc. These risks can arise from individuals that are both inside the organization as well as outside the organization.

[0031] With the growing need to detect and prevent policy violations related to the user, storage, or transmission of sensitive and/or private information, data security and sensitive have become increasingly crucial. To take proactive measures to safeguard sensitive and/or private information, these computing environments often include security infrastructure to enforce access control, data loss prevention, or other processes to secure data from potential vulnerabilities, such as unauthorized access or breaches. One approach to data posture analysis uses scanners (e.g., a computer program running on a physical or virtual machine) that are deployed to access and scan the data stores and detect sensitive and/or private data, or other target data of interest.

[0032] In one approach, scanners are used to scan on-premise data stores and/or data stores in a cloud environment, such as in an organization's cloud accounts, data warehouses, and/or software as a service (SaaS) applications. However, when the data stores being scanned contain large volumes of information, the scanning process can be time-consuming and computationally expensive. Further, that approach is often not scalable as data stores are added and/or grow in size.

[0033] The present technology relates to data discovery in computing environments, such as cloud environments, that provide access to storage resources that store data for users or other endpoints or systems. The described technology can reduce the time and resources required to scan the data stores, while still providing accurate scan results for use in data store classification. Data discovery can identify sensitive and/or private data (or information), which can be used for data posture analysis and management. In one example, the disclosed technology receives a control input identifying a sampling criterion for classifying a data store storing a set of data objects in a computing environment as corresponding to a target data type and deploys one or more scanners configured to select a representative subset of data objects, from the set of data objects, based on the sampling criterion. The representative subset of data objects comprising some, but not all, of

the set of data objects. A scanner result generated by the one or more scanners is received that represents detected instances, in the representative subset of data objects, of one or more pre-defined data patterns of the target data type. A classification result is generated based on a comparison of the number of detected instances of the one or more pre-defined data patterns to a threshold. The classification result represents a classification of the data store as having correspondence to the target data type. A computing action is performed based on the classification result. In one example, the classification of the data store can be utilized to update an infrastructure graph, that graph assets in the cloud environment.

[0034] FIG. 1 is a block diagram illustrating one example of a cloud architecture **100** in which a cloud environment **102** is accessed by one or more actors **104** through a network **106**, such as the Internet or other wide area network. Cloud environment **102** includes one or more cloud services **108-1**, **108-2**, **108-N**, collectively referred to as cloud services **108**. As noted above, cloud services **108** can include cloud storage services such as, but not limited to, AWS, GCP, Microsoft Azure, to name a few.

[0035] Further, cloud services **108-1**, **108-2**, **108-N** can include the same type of cloud service, or can be different types of cloud services, and can be accessed by any of a number of different actors **104**. For example, as illustrated in FIG. 1, actors **104** include users **110**, administrators **112**, developers **114**, organizations **116**, and/or applications **118**. Of course, other actors **120** can access cloud environment **102** as well.

[0036] Architecture **100** includes a cloud security posture analysis system **122** configured to access cloud services **108** to identify and analyze cloud infrastructure and data security posture within cloud environment **102**. Examples of system **122** are discussed in further detail below. Briefly, however, system **122** is configured to access cloud services **108** and identify cloud infrastructure assets, such as by identifying connected resources, entities, actors, etc. within those cloud services, and to identify risks and violations against access to sensitive information. As shown in FIG. 1, system **122** can reside within cloud environment **102** or outside cloud environment **102**, as represented by the dashed box in FIG. 1. Of course, system **122** can be distributed across multiple items inside and/or outside cloud environment **102**.

[0037] Users **110**, administrators **112**, developers **114**, or any other actors **104**, can interact with cloud environment **102** through user interface displays **123** having user interface mechanisms **124**. For example, a user can interact with user interface displays **123** provided on a user device (such as a mobile device, a laptop computer, a desktop computer, etc.) either directly or over network **106**. Cloud environment **102** can include other items **125** as well.

[0038] FIG. 2 is a block diagram illustrating one example of cloud service **108-1**. For the sake of the present discussion, but not by limitation, cloud service **108-1** will be discussed in the context of an account within AWS. Of course, other types of cloud services and providers are within the scope of the present disclosure.

[0039] Cloud service **108-1** includes a plurality of resources **126** and an access management and control system **128** configured to manage and control access to resources **126** by actors **104**. Resources **126** include compute resources **130**, storage resources **132**, and can include other resources **134**. Storage resources **132** can include one or more data stores to be scanned, as discussed below. As used herein, the term “data store” refers to a physical and/or virtual repository for storing data, such as one or more datasets. The data stored in a data store can be located on one or more hardware devices, such as one or more servers, which can be local or remote from one another. In one example, a data store includes one or more database. An example database includes a collection of data that is managed by a database management system (DBMS).

[0040] A virtual data store includes, in one example, a storage system that is abstracted from physical hardware. Therefore, instead of directly accessing physical disks or servers, the virtual data store operates on logical representations of data that can be distributed across various physical devices or locations. Virtual data stores are often implemented in virtualized or cloud environments

where data is managed and accessed through software-defined infrastructure rather than physical hardware. Examples of virtual data stores include virtualized file systems, cloud storage services, and database systems deployed in virtualized environments.

[0041] Compute resources **130** include a plurality of individual compute resources **130-1**, **130-2**, **130-N**, which can be the same and/or different types of compute resources. In the present example, compute resources **130** can include elastic compute resources, such as elastic compute cloud (AWS EC2) resources, AWS Lambda, etc.

[0042] An elastic compute cloud (EC2) is a cloud computing service designed to provide virtual machines called instances, where users can select an instance with a desired amount of computing resources, such as the number and type of CPUs, memory and local storage. An EC2 resource allows users to create and run compute instances on AWS, and can use familiar operating systems like Linux, Windows, etc. Users can select an instance type based on the memory and computing requirements needed for the application or software to be run on the instance.

[0043] AWS Lambda is an event-based service that delivers short-term compute capabilities and is designed to run code without the need to deploy, use or manage virtual machine instances. An example implementation is used by an organization to address specific triggers or events, such as database updates, storage changes or custom events generated from other applications. Such a compute resource can include a server-less, event-driven compute service that allows a user to run code for many different types of applications or backend services without provisioning or managing servers.

[0044] Storage resources **132** are accessible through compute resources **130**, and can include a plurality of storage resources **132-1**, **132-2**, **132-N**, which can be the same and/or different types of storage resources. A storage resource **132** can be defined based on object storage. For example, AWS Simple Storage Service (S3) provides highly-scalable cloud object storage with a simple web service interface. An S3 object can contain both data and metadata, and objects can reside in containers called buckets. Each bucket can be identified by a unique user-specified key or file name. A bucket can be a simple flat folder without a file system hierarchy. A bucket can be viewed as a container (e.g., folder) for objects (e.g., files) stored in the S3 storage resource.

[0045] Compute resources **130** can access or otherwise interact with storage resources **132** through network communication paths based on permissions data **136** and/or access control data **138**. System **128** illustratively includes identity and access management (IAM) functionality that controls access to cloud service **108-1** using entities (e.g., IAM entities) provided by the cloud computing platform.

[0046] Permissions data **136** includes policies **140** and can include other permissions data **142**. Access control data **138** includes identities **144** and can include other access control data **146** as well. Examples of identities **144** include, but are not limited to, users, groups, roles, etc. In AWS, for example, an IAM user is an entity that is created in the AWS service and represents a person or service who uses the IAM user to interact with the cloud service. An IAM user provides the ability to sign into the AWS management console for interactive tasks and to make programmatic requests to AWS services using the API, and includes a name, password, and access keys to be used with the API. Permissions can be granted to the IAM user to make the IAM user a member of a user group with attached permission policies. An IAM user group is a collection of IAM users with specified permissions. Use of IAM groups can make management of permissions easier for those users. An IAM role in AWS is an IAM identity that has specific permissions, and has some similarities to an IAM user in that the IAM role is an AWS identity with permission policies that determine what the identity can and cannot do in AWS. However, instead of being uniquely associated with one person, a role is intended to be assumable by anyone who needs it. Roles can be used to delegate access to users, applications, and/or services that don't normally have access to the AWS resources. Roles can be used by IAM users in a same AWS account and/or in different AWS accounts than the role. Also, roles can be used by compute resources **130**, such as EC2 resources. A service role is a role

assumed by a service to perform actions in an account on behalf of a user. Service roles include permissions required for the service to access the resources needed by the service. Service roles can vary from service to service. A service role for an EC2 instance, for example, is a special type of service role that an application running on an EC2 instance can assume to perform actions.

[0047] Policies **140** can include identity-based policies that are attached to IAM identities can grant permissions to the identity. Policies **140** can also include resource-based policies that are attached to resources **126**. Examples include S3 bucket policies and IAM role trust policies. An example trust policy includes a JSON policy document that defines the principles that are trusted to assume a role. In AWS, a policy is an object that, when associated with an identity or resource, defines permissions of the identity or resource. AWS evaluates these policies when an IAM principal user or a role) makes a request. Permissions in the policy determine whether the request is allowed or denied. Policies are often stored as JSON documents that are attached to the IAM identities (user, groups of users, role).

[0048] A permissions boundary is a managed policy for an IAM identity that defines the maximum permissions that the identity-based policies can grant to an entity, but does not grant the permissions. Further, access control lists (ACLs) control which principles in other accounts can access the resource to which the ACL is attached. ACLs can be similar to resource-based policies. In some implementations of the technology disclosed, the terms “roles” and “policies” are used interchangeably.

[0049] Cloud service **108-1** includes one or more deployed cloud scanners **148**, cloud provider application programming interface(s) (APIs) **150**, one or more data stores **152**, an event log generation component **154**, and can include other items **156** as well. Cloud scanner **148** run locally on the cloud-based services and the server systems, and can utilize elastic compute resources, such as, but not limited to, AWS Lambda resources. Cloud scanner **148** is configured to access and scan the cloud service **108-1** on which the scanner is deployed. Examples are discussed in further detail below. Briefly, however, a scanner accesses the data stored in storage resources **132**, permissions data **136**, and access control data **138** to identify particular data patterns (such as, but not limited to, sensitive string patterns) and traverse or trace network communication paths between pairs of compute resources **130** and storage resources **132**. The results of the scanner can be utilized to identify subject vulnerabilities, such as resources vulnerable to a breach attack, and to construct a cloud attack surface graph or other data structure that depicts propagation of a breach attack along the network communication paths.

[0050] Given a graph of connected resources, such as compute resources **130**, storage resources **132**, etc., entities (e.g., accounts, roles, policies, etc.), and actors (e.g., users, administrators, etc.), risks and violations against access to sensitive information is identified. A directional graph can be built to capture nodes that represent the resources and labels that are assigned for search and retrieval purposes. For example, a label can mark the node as a database or S3 resource, actors as users, administrators, developers, etc. Relationships between the nodes are created using information available from the cloud infrastructure configuration. For example, using the configuration information, system **122** can determine that a resource belongs to a given account and create a relationship between the policy attached to a resource and/or identify the roles that can be taken up by a user.

[0051] Event log generation component **154** is configured to generate event logs **158** (also referred to as activity logs), which record actions taken by a user, a role, or a cloud service as events. Component **154**, in one example, is configured to deploy a database log generator on each of a plurality of databases **162** on storage resources **132**. Databases **162** can store information in any of a plurality of different types of structures. For example, some or all databases **162** can comprise structured query language (SQL) databases that store and organize data in tables with related information. Each table consists of rows and columns, with each row representing a record and each column representing a specific data field. SQL databases allow users to manipulate data

through SQL commands. Some examples of SQL databases include MySQL and PostgreSQL, to name a few. Alternatively, or in addition, some or all databases **162** can include non-SQL databases.

[0052] The respective event log generator, for each respective database **162**, generates a respective database log in event logs **158**. An example log includes log entries that record the database activities, events, etc. Example logs include, but not limited to, audit logs, transaction logs, etc.

[0053] As noted above, in some examples, resources **126** can include AWS EC2 and/or Lambda resources. Also, resources **126** can include AWS Instance Stores and/or AWS Elastic Block Store (EBS) volumes. An EBS volume is a durable, block-level storage device that can attach to a compute instance and used as a physical hard drive.

[0054] Resources **126** can also include an Azure blob identified by a resource URL syntax that assigns each resource a corresponding base URL.

[0055] A cloud storage service or cloud service provider (CSP) can include an organization which hosts services such as networking, software, servers, and/or infrastructure, among others. A CSP can also provide security for the provided services. The services provided by the CSP can relieve a client organization of individual responsibility of setting and managing infrastructure. Examples of CSPs include Amazon Web Services™, Microsoft Azure™, Salesforce™, Google Cloud Platform™, among others.

[0056] Cloud provider APIs **150** are configured to receive calls to access various components in cloud service **108**. For example, cloud provider APIs **150** can access event logs **158** stored in data store **152**. Data store **152** can also store other data items **160** as well.

[0057] A CSP generally provides a number of different interfaces to cloud-computing services, such as a service-provider interface to organizational clients for computing services. A CSP, for example, provides interfaces that allow cloud-computing clients to launch virtual machines, application programs, and other computational entities. A CSP can also provide user interface that allow claims to access, through the Internet, the services provided by the CSP. A client of the CSP can deploy web servers to access, modify, and sending information.

[0058] A cloud account provided by a CSP includes roles that determine user privileges users and what actions can be taken in the cloud account. An identify and access management (IAM) role is managed by the CSP and provides predefined roles that give granular access to specific CSP resources and prevent unwanted access to other CSP resources. For instance, an AWS IAM role includes an AWS identity with a set of permissions policies that each determine what the role can do within an AWS account. An IAM role can be assumed by anyone who needs to require the role.

[0059] FIG. 3 is a block diagram illustrating one example of cloud security posture analysis system **122**. As noted above, system **122** can be deployed in cloud environment **102** and/or access cloud environment **102** through network **106** shown in FIG. 1.

[0060] System **122** includes a cloud account onboarding component **202**, a cloud data scanning and analysis system **206**, a visualization system **208**, and a data store **210**. System **122** can also include a database connection component **212**, one or more processors or servers **214**, and can include other items **215** as well.

[0061] Cloud account onboarding component **202** is configured to onboard cloud services **108** for analysis by system **122**. After onboarding, system **206** is configured to deploy a cloud scanner (e.g., deployed cloud scanner(s) **148** shown in FIG. 2) to the cloud service. In one example, the deployed scanners are on-demand agent-less scanners configured to perform agent-less scanning within the cloud service. One example of an agent-less scanner does not require agents to be installed on each specific device or machine. The scanners operate on resources **126** and access management and control system **128** directly within the cloud service, and generate metadata that is returned to system **122**. Thus, in one example, the actual cloud service data is not required to leave the cloud service for analysis.

[0062] Cloud data scanning and analysis system **206** includes a metadata ingestion component **216**

configured to receive the metadata generated by the deployed cloud scanner(s) **148**. System **206** also includes a query engine **218**, a policy engine **220**, a breach vulnerability evaluation component **222**, one or more application programming interfaces (APIs) **224**, a cloud security issue identification component **226**, a cloud security issue prioritization component **228**, an infrastructure detection and update system **230**, a cloud data scanner deployment and control component **231**, and can include other items **232** as well.

[0063] Query engine **218** is configured to execute queries against the received metadata and generated cloud security issue data. Policy engine **220** can execute security policies against the cloud data and breach vulnerability evaluation component **222** is configured to evaluate potential breach vulnerabilities in the cloud service. APIs **224** are exposed to users, such as administrators, to interact with system **122** to access the cloud security posture data.

[0064] Component **226** is configured to identify cloud security issues and component **228** can prioritize the identified cloud security issues based on any of a number of criteria.

[0065] Visualization system **208** is configured to generate visualizations of the cloud security posture from system **206**. Illustratively, system **208** includes a user interface component **234** configured to generate a user interface for a user, such as an administrator. In the illustrated example, component **234** includes a web interface generator **236** configured to generate web interfaces that can be displayed in a web browser on a client device.

[0066] Visualization system **208** also includes an infrastructure graph generator component **238**, a cloud attack surface graph generator component **240**, and can include other items **242** as well. Infrastructure graph generator component **238** is configured to generate a graph or other representation of the relationships between resources **126** or other infrastructure assets such as privileges, roles, etc. For example, component **238** can generate a cloud infrastructure graph, such as a map, that graphically depicts pairs of compute resources and storage resources as nodes and network communication paths as edges between the nodes.

[0067] Cloud attack surface graph generator component **240** is configured to generate a surface graph or other representation of vulnerabilities of resources to a breach attack. In one example, the representation of vulnerabilities can include a cloud attack surface map that graphically depicts propagation of a breach attack along network communication paths as edges between nodes that represent the corresponding resources.

[0068] Data store **210** stores metadata **244** obtained by metadata ingestion component **216**, sensitive data profiles **246**, detected event logs **248**, and can store other items **250** as well. Examples of sensitive data profiles **246** are discussed in further detail below. Briefly, however, sensitive data profiles **246** can identify target data patterns that are to be categorized as sensitive or conforming to a predefined pattern of interest. Sensitive data profiles **246** can be used as training data for data classification. Examples of data classification are discussed in further detail below. For instance, however, pattern matching can be performed based on the target data profiles. Illustratively, pattern matching can be performed to identify instances of data patterns corresponding to social security numbers, credit card numbers, other personal data, medical information, to name a few. In one example, artificial intelligence (AI) is utilized to perform named entity recognition (e.g., natural language processing modules can identify sensitive data, in various languages, representing names, company names, locations, etc.).

[0069] Database connection component **212** is configured to connect to, or access, databases in the data stores of the resources being analyzed by system **122**. Examples are discussed in further detail below. Briefly, however, database connection component **212** can receive user access credentials, such as a username and password, for each database of a plurality of databases to be accessed in the cloud environment and scanned by the deployed scanners. In another example, database connection component **212** can be configured to connect to representations of the databases that are accessed using a different authentication requirement, than the databases themselves. For example, database connection component **212** can identify and connect to snapshots of the databases through cloud

provider APIs **150**, discussed above. Database connection component **212** can access event logs **158**, for the identified databases, from data store **152**.

[0070] Infrastructure detection and update system **230** includes a trigger detection component **252**, a scan criterion generation component **254**, an event log detection system **256**, a graph update component **258**, a post-processing action generator component **260**, an output component **262**, and can include other items **264** as well.

[0071] Operation of system **230** is discussed in further detail below. Briefly, however, trigger detection component **252** is configured to detect a triggering criterion corresponding to initiation of an update scan of the infrastructure posture of cloud environment **102**. The update scan is configured to identify changes to the infrastructure posture that occurred subsequent to a previous scan and update performed by system **230**. In this way, the trigger criterion can be based on a temporal criterion, such as the duration and/or end time of a prior scan.

[0072] Scan criterion generation component **254** is configured to generate one or more scan criterion, for example based on the temporal triggering criterion. Event log detection system **256** includes an orchestration engine configuration component **266** and an asset discovery component **268**. System **256** can include other items as well.

[0073] Orchestration engine configuration component **266** configures an orchestration engine to deploy one or more micro services to analyze event logs (herein referred to as log analyzer microservices) and implement data scanner(s). For example, the event logs can be a plurality of database activity logs for a plurality of databases in cloud environment **102**. Here, the orchestration engine deploys, in parallel, a plurality of microservices to analyze those logs.

[0074] An example orchestration engine includes a tool or platform that automates and manages the deployment, configuration, coordination, and monitoring of complex applications or services. The orchestration engine can provide a centralized control system that streamlines the management of multiple components and resources, such as containerized applications, across distributed systems. In addition, the orchestration engine can include features such as load balancing, automatic failover, and resource allocation to help ensure that the application is highly available and can scale dynamically to handle changing workloads. In one example, each microservice is configured to perform a specific function that is deployable and scalable independently of other services. The microservice(s) return analysis results to the orchestration engine, such as in the form of metadata from the event logs and/or metadata indicative of a scan of data objects in the data store.

[0075] Based on the results, graph update component **258** is configured to update the infrastructure graph, such as by adding nodes, adding edges, deleting nodes, deleting edges, or editing/changing existing nodes and/or edges in the graph. The updates can reflect, for example, changes to infrastructure assets, such as events that add, delete or change resources, such as compute resources (e.g., EC2 instances), storage resources (e.g., S3 instances), roles in the cloud environment, privileges, to name a few.

[0076] Output component **262** can output the analysis results to other systems. In one example, alerts or user interface displays, such as reports, can be provided to an administrator or other user.

[0077] Remedial action generator component **270** is configured to generate a remedial action that can be output to an administrator or user as a suggested action. Alternatively, or in addition, component **270** can generate control signals that perform the remedial action within the cloud environment. Examples are discussed in further detail below.

[0078] FIG. **4** is a block diagram illustrating one example of a deployed scanner **148**. Scanner **148** includes a resource identification component **402**, a permissions data identification component **404**, an access control data identification component **406**, a cloud infrastructure scanning component **408**, a cloud data scanning component **410**, an output component **412**, a log analyzer microservice component **414**, and can include other items **416** as well. FIG. **4** also illustrates that some or all components of and/or functionality performed by system **230** can be on or otherwise associated with deployed scanner **148**.

[0079] Resource identification component **402** is configured to identify the resources **126** within cloud service **108-1** (and/or other cloud services **108**) and to generate corresponding metadata that identifies these resources. Permissions data identification component **404** identifies the permissions data **136** and access control data identification component **406** identifies access control data **138**. Cloud infrastructure scanning component **408** scans the infrastructure of cloud service **108** to identify the relationships between resources **130** and **132** and cloud data scanning component **410** scans the actual data stored in storage resources **132**. Output component **412** is configured to output the generated metadata and content-based classification results to cloud security posture analysis system **122**.

[0080] The metadata generated by scanner **148** can indicate a structure of schema objects in a data store. For example, where the schema objects comprise columns in a data store having a tabular format, the returned metadata can include column names from those columns. A content-based data item classifier is configured to classify data items within the schema objects, based on content of those data items. Examples are discussed in further detail below.

[0081] Cloud data scanning component **410** includes a sampling component **418**, a content-based classifier **420**, a classification metadata generator component **422**, and can include other items **424** as well. Sampling component **418** is configured to define or otherwise obtain a sampling criterion that is utilized for sampling a data store to be scanned. Using the sampling criterion, scanner **148** selects a representative subset of data objects in the data store to be utilized by scanner **148** in performing content-based classification. Operation of sampling component **418** is discussed in further detail below.

[0082] Content-based classifier **420** is configured to perform content-based classification to classify data objects (e.g., data files, etc.) in the data store. Content-based classifier **420** includes a data object accessing component **426**, a target data entity comparison component **428**, a result generator **430**, a confidence score generator **432**, and can include other items **434** as well. Data object accessing component **426** is configured to access the data objects.

[0083] Content-based classifier **420** can perform content-based classification using target data entity comparison component by comparing the data objects to predefined target data profiles (e.g., sensitive data profiles). The content-based classifier can be any of a wide variety of classifiers including, but not limited to, a rules-based classifier, a machine learned classifier (e.g., a deep learning system), a heuristics-based classifier, or other type of classifier or classification model.

[0084] An example classifier includes a classification engine that evaluates extracted content, extracted from the data objects. For example, a portion of content can be matched to a content rule if a characteristic of the content satisfies a condition of the content rule. For instance, a classifier can compare the extracted content to the target data profiles using one or more similarity measures. A similarity measure can be used by the classifier to determine whether a data item extracted from a content object is determined to match a given target data profile. Examples of similarity measures include, but are not limited to, k-gram overlap, edit distance, Cosine similarity, Manhattan distance, Euclidean distance, Levenshtein distance, to name a few.

[0085] Log analyzer microservice component **414** is configured to analyze one or more event logs to identify incremental changes to infrastructure assets. Component **414** includes a log analyzer microservice that is employed and managed by an orchestration engine, as discussed above. Component **414** includes a log parser component **438**, a log query component **440**, a result output component **442**, and can include other items **444** as well.

[0086] Log parser component **438** is configured to parse event logs, such as database event logs in a given database, on which the scanner is deployed. The parsing can be done in any of a number of ways. In one example, log parser component **438** accesses each of a number of entries in the event log and identifies constituent data fields that identify various characteristics or parameters from a respective event or activity in the cloud environment. For example, the parsing can include an identification of what service (e.g., a resource or other asset) had a change, a time that the change

occurred, what geographic region the change occurred in, what entity was involved in the change (e.g., a requestor identity), as well as any other attributes. These attributes can be returned as metadata by result output component **442**.

[0087] FIG. **5** is a flow diagram **500** showing an example operation of system **122** in on-boarding a cloud account and deploying one or more scanners. At block **502**, a request to on-board a cloud service to cloud security posture analysis system **122** is received. For example, an administrator can submit a request to on-board cloud service **108-1**.

[0088] At block **504**, an on-boarding user interface display is generated. At block **505**, user input is received that defines a new cloud account to be on-boarded. The user input can define a cloud provider identification **506**, a cloud account identification **508**, a cloud account name **510**, access credentials to the cloud account **512**, and can include other input **514** defining the cloud account to be on-boarded.

[0089] At block **516**, the cloud account is authorized using roles. For example, administrator access (block **518**) can be defined for the cloud scanner using IAM roles. One or more cloud scanners are defined at block **520** and can include, but are not limited to, cloud infrastructure scanners **522**, cloud data scanners **524**, vulnerability scanners **526**, event log scanners **528**, or other scanners **530**.

[0090] At block **532**, the cloud scanners are deployed to run locally on the cloud service, such as illustrated in FIG. **2**. The cloud scanners discover resources at block **534**, scan data in the resources at block **536**, event logs at block **538**, and can find vulnerabilities at block **540**. As discussed in further detail below, a vulnerability can be identified based on finding a predefined risk signature in the cloud service resources. The risk signatures can be queried upon, and define expected behavior within the cloud service and locate anomalies based on this data.

[0091] At block **542**, if more cloud services are to be on-boarded, operation returns to block **504**. At block **544**, the scan results from the deployed scanners are received. As noted above, the scan results include metadata (block **546**) and/or data item classifications (block **548**) generated by the scanners running locally on the cloud service.

[0092] At block **550**, one or more actions are performed based on the scan results. At block **552**, the action includes security issue detection. For example, a breach risk on a particular resource (such as a storage resource storing sensitive data) is identified. At block **554**, security issue prioritization can be performed to prioritize the detected security issues. Examples of security issue detection and prioritization are discussed in further detail below. Briefly, security issues can be detected by executing a query against the scan results using vulnerability or risk signatures.

[0093] The risk signatures identify criterion such as accessibility of the resources, access and/or permissions between resources, and data types in accessed data stores. Further, each risk signature can be scored and prioritized based on impact. For example, a risk signature can include weights indicative of likelihood of occurrence of a breach and impact if the breach occurs.

[0094] A risk signature can be defined based upon any of a wide variety of criteria. For example, a risk signature can identify one or more configurations or settings of compute resources **130**. Examples include, but are not limited to, a configuration that indicates whether the compute resource provides accessibility to a particular type of data, such as confidential data, medical data, financial data, personal data, or any other type of private and/or sensitive content. In another example, a risk signature indicates that a compute resource is publicly accessible, includes a public Internet protocol (IP) address, or has IP forwarding enabled. In another example, a risk signature indicates that a compute resource has monitoring disabled, has no IAM role assigned to the compute resource, has backup disabled, data encryption disabled, and/or a low or short backup retention policy. Also, a risk signature can identify password policies set for the compute resource. For instance, a risk signature can indicate a lack of minimum password policies, such as no minimum password length, no requirement of symbols, lowercase letters, uppercase letters, numbers, or password reuse policy. Also, a risk criterion can indicate a location of the compute resource, such as whether the compute resource is located outside of a particular region.

[0095] Risk signatures can also indicate configurations and/or settings of storage resources **132**. For example, the configurations and settings can indicate authentication or permissions enforced by the storage resource, such as whether authentication is required for read, write, delete, synchronization, or any other operation. Also, the risk signature can indicate whether multi-factor authentication is disabled for the storage resource, as well as a breadth of permissions grants (e.g., whether all authenticated users are granted permissions within the storage resource). Also, a risk signature can indicate whether encryption is enabled by default, a password policy enforced by the storage resource, whether the storage resource is anonymously accessible, publicly accessible, has a key management service disabled, has logging disabled, life cycle management disabled, whether the storage resource is utilized for website hosting, has geo-restriction disabled, or has backup functionality disabled. Also, the risk signature can indicate a type of data stored by the storage resource, such as the examples discussed above.

[0096] The action can further include providing user interfaces at block **556** that indicate the scan status (block **558**), a cloud infrastructure representation (such as a map or graph) (block **560**), and/or a cloud attack surface representation (map or graph) (block **562**). The cloud attack surface representation can visualize vulnerabilities.

[0097] Remedial actions can be taken at block **564**, such as creating a ticket (block **566**) for a developer or other user to address the security issues. Of course, other actions can be taken at block **568**. For instance, the system can make adjustments to cloud account settings/configurations to address/remedy the security issues.

[0098] FIG. **6** illustrates one example of a user interface display **600**, that can be displayed at block **552**. Display **600** provides a dashboard for a user which provides an overview of on-boarded cloud service accounts. The dashboard identifies a number of users **602**, a number of assets **604**, a number of data stores **606**, and a number of accounts **608**. A data sensitivity pane **610** includes a display element **612** that identifies a number of the data stores that include sensitive data, a display element **613** that identifies a number of users with access to the sensitive data, a display element **614** that identifies a number of resources having sensitive data, and a display element **616** that identifies a number of risks on the data stores having sensitive data. Further, graphs or charts can be generated to identify those risks based on factors such as status (display element **618**) or impact (display element **620**).

[0099] Display element **620** illustratively categorizes the risks based on impact as well as the likelihood of occurrence of those risks. Risk categorization is discussed in further detail below. Briefly, however, display element **620** stratifies one or more of breach likelihood scores or breach impact scores categories representing different levels of severity, such as high, medium, and low severity levels. In one example, display element **620** is color coded based on the degree of impact of the risk (e.g., high impact is highlighted in red, medium impact is highlighted in yellow, and low impact is highlighted in green).

[0100] FIG. **7** is a flow diagram **700** illustrating one example of cloud infrastructure scanning performed by cloud scanner **148** deployed in cloud service **108-1**. At block **702**, an agent-less scanner is executed on the cloud service. For example, a compute resource can utilize a serverless function to perform a scan on the cloud service.

[0101] At block **704**, the scanner discovers the compute resources **130** and, at block **706**, the storage resources **132**. Sensitive data can be discovered at block **708**. The agent-less scanner does not require a proxy or agent running in the cloud service, and can utilize server-less containers and resources to scan the documents and detect sensitive data. The data can be accessed using APIs associated with the scanners. The sensitive data can be identified using pattern matching, such as by querying the data using predefined risk signatures.

[0102] At block **710**, access paths between the resources are discovered based on permissions data **136** (block **712**), and/or access control data **138** (block **714**). A rule processing engine, such as using JSON metadata, can be utilized to analyze the roles and policies, and can build access

relationships between the nodes representing the resources. The policies can be decoded to get access type (allow, deny, etc.) and the policy can be placed in a node to link from a source to target node and create the access relationship. At block **716**, metadata indicative of the scanning results is generated and outputted by metadata output component **272**.

[0103] FIGS. **8-1** and **8-2** (collectively referred to as FIG. **8**) provide a flow diagram **800** illustrating one example of infrastructure analysis and query execution. At block **802**, permissions data and access control data for pairs of compute and storage resources is accessed. Policy data is accessed at block **804**. For example, the policy data can include identity-based policies (block **806**), resource-based policies (block **808**), permissions boundaries (block **810**), service control policies (SCP) (block **812**), session policies (block **814**) as well as other policies (block **816**).

[0104] At block **818**, network communication paths are traced between the pairs of resources. Tracing the network communication path can be based on the permissions data at block **820**, the access control data at block **822**, the policy data at block **824**, and/or other data at block **826**.

[0105] At block **828**, a cloud infrastructure map is constructed. An example of a cloud infrastructure map includes nodes that graphically represent pairs of compute and storage resources (block **830**), and edges that represent network communication paths between the resources (block **832**). At block **834**, the map graphically depicts metadata associated with the pairs of resources. For example, a graphical metadata depiction is expandable or collapsible via user selection, as represented at block **836**. The metadata can be grouped across metadata categories at block **838**, such as based on cloud-sourced metadata at block **840**, derived metadata at block **842**, locally annotated metadata at block **844**, or based on other metadata categories at block **846**.

[0106] The cloud infrastructure map can also graphically depict anomalous configured access instances at block **848**. For example, block **848** can detect different levels of access among resources that connect to a common network component, as represented at block **850**. At block **852**, the map graphically depicts anomalous actual access instances in the cloud environment. For instance, the instances can be detected from access logs at block **854**. User annotated tags for the resources can be depicted in the map at block **856** as well.

[0107] At block **858**, a query is received. The query can include a search term, a content category, a data privacy policy, a temporal period, and can include other items as well.

[0108] Alternatively, or in addition, a filter criterion is received can be received at block **860**. The filter criterion can be based on the metadata, based on applications running on at least one pair of resources, and/or based on one or more networks in the cloud environment.

[0109] The query and/or filter criterion are executed at block **862** and results are returned at block **864**. For example, the query results can identify a subset of the pairs of resources that contain searched content.

[0110] FIG. **9** illustrates a user interface display **900** that includes a visualization of cloud infrastructure assets and relationships between those assets, such as access communication paths. The visualization in FIG. **9** can be rendered as a cloud infrastructure graph (e.g., map) that shows relationships between compute and storage resources and/or mappings between users, roles, and resources, based on the permissions data and the access control data. Further, the visualization can be augmented using sensitivity classification data to represent propagation of breach attack along communication paths.

[0111] As shown in FIG. **9**, nodes **902** represent compute resources and nodes **904** represent storage resources. Illustratively, the storage resources include data stores or buckets within a particular cloud service. Nodes **906** represent roles and/or users. The links (e.g., access paths) or edges **908** between nodes **902** and **906** represent that compute resources that can access the particular roles represented by nodes **906**. The edges or links **910** represent the storage resources that can be accessed by the particular roles or users represented by nodes **906**.

[0112] Based on these relationships between compute and storage relationships, display elements can be rendered along, or otherwise visually associated with, the edges **908** and/or **910**, to identify

and graphically depict the propagation of breach attack. For instance, vulnerability display elements can be rendered in association with edges **908** and/or **910** to identify that a subject vulnerability signature (e.g., one or more risk signatures shown in FIG. **9**) has been identified in the data, based on querying the permissions and access control data using the subject vulnerability signature. For example, display element **909** represents a risk signature between nodes **903** and **912** and display element **911** represents (such as by including a description, icon, label, etc.) a risk signature between nodes **912** and **922**. Each display element **909**, **911** can represent (such as by including a description, icon, label, etc.) corresponding likelihood and impact scores, can be actuatable to render details of the subject vulnerability, such as in a display pane on display **900**. The details can include which risk signature has been matched, which sensitive data is at risk, etc. [0113] FIGS. **10-1** and **10-2** (collectively referred to as FIG. **10**) provide a flow diagram **1000** of an example operation of infrastructure posture analysis using incremental change detection. For sake of illustration, but not by limitation, FIG. **10** will be discussed in the context of system **200** discussed above with respect to FIG. **3**.

[0114] At block **1002**, an initial cloud infrastructure graph is obtained. An example infrastructure graph includes nodes that represent resources in the cloud environment and edges, between the nodes, that represent relationships between the resources. Examples are discussed above with respect to FIGS. **8** and **9**.

[0115] At block **1004**, a triggering criterion for initiation of an update scan is detected. For example, the detection can be automatic at block **1006**, manual at block **1008**, or another type of input, as represented at block **1010**.

[0116] In one example, the triggering of the update scan is performed automatically according to a periodic schedule, as represented at block **1012**. For example, the update scan can be performed every x number of hours (e.g., one hour, two hours, three hours, etc.). In another example, the period schedule can define that a next update scan is to occur x number of minutes (e.g., fifteen minutes, etc.) after the previous scan has completed. In this way, the periodic schedule can be dynamic based on a duration of a previous update scan, as represented at block **1014**.

[0117] An incremental change detector is invoked. One example of an incremental change detector is discussed below with respect to FIG. **11**. Briefly, however, the incremental change detector can receive an input parameter at block **1018**, such as a target time within which to detect changes in the cloud infrastructure. For example, the input parameter indicates a time at which a previous scan ended. In this way, the incremental change detector will identify any changes that occur since the last time the cloud infrastructure was scanned to update the infrastructure graph.

[0118] At block **1020**, the cloud environment is scanned by the incremental change detector. The incremental change detector can invoked in other ways as well, as represented at block **1022**.

[0119] At block **1024**, scan results are received from the change detector. The scan results can include metadata received at block **1026**, indicative of changes to one or more infrastructure assets. A change to an infrastructure asset can include, but are not limited to, the addition of a new asset (block **1028**), the deletion of an existing asset (block **1030**), the changing or editing of an existing asset (block **1032**), or other types of changes (block **1034**).

[0120] Further, the asset can include, but is not limited to, a particular service or service configuration, as represented at block **1036**. An asset can include a compute resource (block **1038**), a storage resource (block **1040**), a privilege (block **1042**), and/or a role (block **1044**).

[0121] For sake of illustration, an example change indicates that a user y launched a new compute resource (e.g. a particular EC2 instance) of type z at a particular time t. The returned metadata can indicate the user, a unique identifier that identifies the launched resource, and the type of the resource. The metadata can also indicate the roles and/or privileges given to that resource.

[0122] At block **1046**, a set of infrastructure asset changes are identified based on the scan results. Based on the set of infrastructure asset changes, post processing steps can be constrained to the set of infrastructure asset changes. In this way, the post processing is not performed for the entire

infrastructure posture, but the post processing is narrowly tailored to only the particular incremental changes detected during the current scan.

[0123] For example, at block **1048**, a particular change in the set of infrastructure asset changes is selected and, at block **1050**, the extracted metadata for that change is obtained.

[0124] The metadata extracted at block **1050** can include metadata identifying which infrastructure asset (e.g., the target of the event such as a particular storage resource) which is represented at block **1052**. Alternatively, or in addition, the extracted metadata can identify when the event occurred (block **1054**), which region the event occurred in (block **1056**), which entity was involved in the change (block **1058**) and can include other metadata (block **1060**) as well. The entity involved in the change at block **1058** can include the identity of a particular requestor (a user, role, service, etc.).

[0125] At block **1062**, an asset-specific scan and/or query is generated for additional details of the particular change. For example, an asset-specific API call can be generated to access the particular infrastructure asset that was the target of the particular change being analyzed. The call can be utilized, for example, to perform a data scan on the data in the resource, as represented at block **1064**. For example, the data scan can identify which data profiles reside in the resource, such as sensitive data. Of course, the additional details for the particular change can be identified in other ways as well, as represented at block **1066**.

[0126] At block **1068**, the operation determines whether there are more changes in the set of infrastructure asset changes to be analyzed. If so, operation returns to block **1048** in which another change from the set is selected to obtain the extracted metadata and generate a scan or query for additional details on a particular change.

[0127] At block **1070**, the infrastructure graph can be updated in a post processing step. An example of performing operations in an infrastructure graph is discussed above with respect to FIG. **8**. Block **1070** can include similar features as those described with respect to FIG. **8**.

[0128] Updating the infrastructure graph can include, for example, adding nodes at block **1072**, adding edges at block **1074**, deleting nodes at block **1076**, deleting edges at block **1078**, or other updates at block **1080**. For example, the configuration represented by nodes and/or edges can be modified to represent changes in the infrastructure posture.

[0129] For sake of illustration, but not by limitation, if one of the events in the set of infrastructure asset changes identifies that a user added a new compute resource that is accessible by a compute resource with a particular role, a new node and edge can be added to the infrastructure graph to identify the new resource and relationship to the existing resource(s) identified in the graph.

[0130] At block **1082**, other post processing steps can be applied based on the changes as well. For example, the post processing steps can include filtering a set of risk signatures to obtain a filtered set of risk signatures, and applying the filtered set of risk signatures at block **1084**. Examples of block **1084** are discussed below with respect to FIG. **12**. Briefly, however, block **1084** can include identifying attributes in the set of infrastructure asset changes, and determining which risk signatures are applicable to those assets. Block **1084** can include, but is not limited to, utilizing tags that identify the type of assets (whether the assets being changed are ECs instances, S3 buckets, RDS assets, etc.). Using this information, a set of risk signatures can be filtered into a reduced, filtered set of risk signatures, and only those risk signatures in the reduced set can be applied. Filtering the risk signatures before application to the infrastructure graph reduces processing time and load required to process the graph for matching risk signatures. Of course, other post processing steps can be performed at block **1086**.

[0131] At block **1088**, operation determines whether to continue by, for example, initiating a subsequent update scan based on triggering criteria detected at block **1004**.

[0132] FIG. **11** is a flow diagram **1100** illustrating one example of invoking an incremental change detector at block **1016**.

[0133] At block **1102**, input parameters are received. As noted above, the input parameters can

include, but are not limited to, a temporal criteria at block **1104**. For example, the temporal criteria can indicate a time that a previous scan ended and thus a beginning time instance for a time window within which to perform the incremental change detection. The input parameters can also indicate a type to detect changes, as represented at block **1106**. Of course, other input parameters can be received as well, as represented at block **1108**.

[0134] At block **1110**, a plurality of target constructs are identified for which to perform log scanning. In one example, block **1110** includes discovering a plurality of databased in the cloud environment, such as by using cloud provider APIs at block **1112**, and/or identifying databases having logging enabled at block **1114**. Of course, the locations to scan for change detection can be performed in other ways as well, as represented at block **1116**.

[0135] At block **1118**, an orchestration engine is configured that deploys and manages a plurality of log analyzer microservices to operate on respective event logs. In one example, block **1118** includes operating on a plurality of respective event logs, in parallel, that correspond to a plurality of databases. The orchestration engine and/or log analyzer microservices can include compute instances and/or virtual machines at block **1120**. For example, a serverless compute instance, such as a lambda instance, can be utilized. Of course, the plurality of log analyzer microservices can be deployed and managed in other ways as well, as represented at block **1122**.

[0136] At block **1124**, each log analyzer microservice runs to scan the corresponding event log(s) based on the input parameters received at block **1102**. Block **1124** can include identifying the asset change events at block **1126**, that match the input parameters. Examples of event logs that can be scanned include, but are not limited to, audit logs at block **1128**, transaction logs at block **1130**, or other types of logs at block **1132**.

[0137] At block **1134**, scan results are received from the log analyzer microservices. Block **1134** can include results that identify write events (block **1136**) identified by the log analyzer microservices in the event logs. A write event includes activities such as creating a new resource, changing an existing resource, deleting a resource, etc. This is contrasted to a read event in which the activity includes only a data access.

[0138] At block **1138**, the event log can be filtered based on the input parameters. In this way, the scan results received from the log analyzer microservice include filtered event logs received as extracted metadata at block **1140**. Alternatively, or in addition, the scan result can include asset tags (block **1142**), that provide information (such as resource type, etc.) of the assets that were changed. Of course, the scan results can be received in other ways as well, as represented at block **1144**. At block **1146**, analysis results are outputting based on the scan results.

[0139] FIG. **12** is a flow diagram **1200** illustrating one example of filtering and applying risk signatures at block **1084**. As represented at block **1204**, the tags correspond to asset tags, such as identifying the types of resources to which a given risk signature is relevant.

[0140] At block **1202**, a tag set of risk signatures is accessed. FIG. **13** illustrates one example of a user interface display **1300**, that renders a set of tagged risk signatures. As shown in FIG. **13**, display **1300** provides a searchable interface to access, modify, add, and/or delete risk signatures that can be applied to the cloud infrastructure posture. Each risk signature defines a set of criteria that the resources and data in the cloud environment can be queries upon to identify indications of vulnerabilities in the cloud service.

[0141] The example of FIG. **13** provides the set of tagged risk signatures in a tabular format, having a plurality of rows **1302** each representing a different risk signature. Each row includes a plurality of data fields or cells arranged in a plurality of columns **1304**. The plurality of columns includes a name column **1306**, a signature ID column **1308**, a description column **1310**, a result header column **1312**, a resource column **1314**, a tags column **1316**, a breach likelihood factor column **1318**, and a breach impact factor column **1320**. Column **1306** includes a name of the given risk signature, uniquely identified by a signature ID in column **1308**. The description column **1310** describes the risk signature, and the result header column **1312** identifies the information to be

returned in results that match the risk signature. The tags column **1316** identifies one or more tags, for the given risk signature. The tags identify characteristics of the resources to which the risk signature applies. For example, the tags **1322** for risk signature **1324** indicate that the risk signature is applicable to S3 resources. This, of course, is for sake of example only.

[0142] Likelihood factor column **1318** indicates a likelihood factor that is assigned to the risk signature and an impact factor column **1320** indicates an impact factor assigned to the risk signature. The likelihood factor indicates a likelihood assigned to occurrence of the risk signature and the impact factor assigns an impact to the cloud service assigned to the occurrence of the risk signature. For sake of illustration, a likelihood factor of ten (out of a scale of ten) indicates that the vulnerability is likely to occur if the risk signature is identified in the cloud posture data, whereas a likelihood factor of one indicates a low likelihood. Similarly, an impact factor of ten (out of a scale of ten) indicates that the vulnerability is considered to have a high impact, whereas an impact factor of one indicates the vulnerability is considered to have a low impact on the cloud service.

[0143] Referring again to FIG. **12**, at block **1206**, a set of tags from the incremental scan of infrastructure resources is collected. For sake of example, assume that the operation described in FIG. **10** returns fifty asset changes that have occurred in the cloud infrastructure. Block **1206** analyzes these changes to identify the tags relevant to the resources involved in the changes. For instance, if the set of changes involved EC2 instances and S3 buckets, block **1206** returns the tags “S3” and “EC2.”

[0144] At block **1208**, the risk signatures are filtered to obtain a filtered set. To illustrated, in the above example where two tags “EC2” and “S3” are returned, block **1208** returns the risk signatures in the set of tagged risk signatures accessed at block **1202** that include at least one of the tags “S3” or “EC2.” In other words, block **1208** operates to filter the collection of risk signatures to identify only those risk signatures that are tagged as applicable to the type of asset(s) involved in the changes.

[0145] Thus, when the risk signatures are applied at block **1210**, only the filtered set of risk signatures is applied to the infrastructure graph, which reduces the processing expense and reduces processing latency in finding instances of matching risk signatures in the infrastructure graph.

[0146] At block **1212**, results of the instances of matching risk signatures are outputted, for example in a user interface display. One example display is illustrated in FIG. **14**.

[0147] As shown in FIG. **14**, display **1400** includes a plurality of rows **1402** that indicate matching instances of risk signatures. Each row includes a description field in column **1404**, a resource name field in column **1406**, an account ID field in column **1408**, a tags field in column **1410**, an impact field in column **1412**, a likelihood field in column **1414**, and a signature ID field in column **1416**. The signature ID field in column **416** identifies the signature ID that was matched by the instance identified in the row. The tags field identifies which tags were matched between the infrastructure change and the risk signature that was matched for the given row.

[0148] FIG. **15** illustrates one example of cloud data scanner deployment and control component **231**. Component **231** includes a sampling criterion component **1502**, a user interface component **1504**, a classifier training component **1506**, a metadata receiving component **1508**, a data store **1510**, and can include other items **1512** as well. Data store **1510** can store classifier training data **1514** and can store other items **1516** as well.

[0149] Training data **1514** illustratively includes target data profiles (also referred to as target data profiles or entities) that represent a collection of different data types. For example, training data **1514** can include sensitive data profiles, which can represent sensitive data types, patterns, and/or signatures. The sensitive data types can include personally identifiable information (PII), protected health information (PHI), payment card information (PCI), financial information, tax information, or other types of sensitive data types. Examples include various types of financial data, such as, but not limited to, credit card numbers, bank account numbers, etc. Also, the training data **1514** can identify personal information, such as social security numbers, phone numbers, email addresses,

etc.
[0150] For sake of illustration, but not by limitation, an example target data profile for an internet protocol (IP) address indicates a form “x.x.x.x”, where x is an octet and must be a decimal value between zero and two hundred and fifty five. As another example, a target data profile for a credit card number indicates a string of fifteen or sixteen characters that begins with a three, four, or five.
[0151] Classifier training component **1506** is configured to train classifier **420** using training data **1514**.

[0152] User interface component **1504** is configured to generate a user interface with user interface display elements and controls to configure the scanner to be deployed. For example, this can include input mechanisms that facilitate user selection of the sampling criterion used in scanning a data store.

[0153] FIG. **16** is a flow diagram **1600** illustrating one example of configuring and deploying a data scanner on a data store using one or more sampling criterion. At block **1602**, a request is received to configure one or more data store scanners. For example, the request can include receiving an on-boarding request to on-board a new cloud account, to add a new data to a cloud account, etc.

[0154] At block **1604**, a target data type is identified to scan in the one or more data stores. The target data type can include, for example, a particular data pattern that represents sensitive and/or private information, such as the examples described above (e.g., PII data, PHI data, PCI data, financial information, tax information, etc.).

[0155] At block **1606**, a control input identifying one or more sampling criterion is received. In one example, a user interface is generated and includes one or more user interface controls that are actuatable by a user to select a particular data store (block **1608**), a user interface control that is actuatable to select a sampling criterion (block **1610**), and/or other user interface elements (block **1612**). Examples of user interface controls include, but are not limited to, checkboxes, radio buttons, dropdown lists, list boxes, buttons, toggles, text fields, date fields, breadcrumbs, sliders, search fields, paginations, sliders, tags, icons, etc.

[0156] The data store can include unstructured data stores (block **1614**), structured data stores (block **1616**), or other data stores (block **1618**). Examples of unstructured data stores include, but are not limited to, S3 buckets, GCP blob containers, azure blob containers, etc.

[0157] The user interface control at block **1610** is actuatable to select one or more sampling criterion to be used in configuring the data scanner. Examples of the sampling criterion selected at block **1610** includes a random sampling criterion at block **1620**. The random sampling criterion is used by the data scanner to randomly select a number of data objects (such as data files) to be scanned within the data store based on the total number of data objects in that data store. For example, the number of data objects to be sampled can be a defined percentage of the total number of data objects in the data store.

[0158] The number of data objects to be sampled can be defined based on a target (e.g., maximum acceptable) target error margin, as represented at block **1622**. The target error margin can be selected through the user interface, or it can be identified programmatically, e.g., based on the size of the data store. Based on the selected target error margin, the number of data objects to sample from the data store can be identified through a lookup table. One example of a lookup table is shown below in Table 1.

TABLE-US-00001		TABLE 1		Acceptable margin of		Size of population error		Large		5000		2500		1000										
500	200	±20%	24	24	24	23	23	22	±15%	43	42	42	41	39	35	±10%	96	94	93	88	81	65	±7.5%	171
165	160	146	127	92	±5%	384	357	333	278	217	132	±3%	1067	880	748	516	341	169						

[0159] Through a user input mechanism, a user can select the target error margin for the scan, and the operation can determine the number of data objects to selected based on the size (e.g., total number of data objects) of the data store.

[0160] Another example of a sampling criterion is based on a hierarchical structure of data objects

in the data store, as represented at block **1624**. For example, in a data store, files are stored in a directory structure. The sampling criterion at block **1624** selects data objects to scan based on which directory those data objects reside on in the data store. For example, the sampling criterion can select the data objects to scan such that the scanned data objects are evenly distributed (or some other defined distribution) amongst all the directories in the data store. Alternatively, or in addition, the sampling criterion can define that at least one data object is to be scanned in each directory. This, of course, is for sake of example only.

[0161] Another example of a sampling criterion is a metadata-based sampling criterion (block **1626**). The metadata-based sampling criterion selects data objects to sample based on metadata associated with those data objects. For example, the metadata can define object types (e.g., file types), tags and/or data object sizes (e.g., file sizes). For example, data objects can be ignored if they include metadata indicating similar characteristics as scanned data objects. In another example, the metadata-based sampling criterion selects data object such that there is at least one data object of each given object type that is selected.

[0162] Another example of a sampling criterion can include a timestamp-based sampling criterion (block **1628**). The timestamp-based sampling criterion is configured to select the data objects to scan based on when each data object was last modified, as identified by its timestamp. In one example, the most recently modified data objects are selected first, for sampling.

[0163] Another example of a sampling criterion includes an exclusion-based sampling criterion (block **1630**). The exclusion-based sampling criterion defines parameters for excluding certain data objects from the scan. For example, a user can select particular data object types, storage locations, directories, etc., to be ignored in the scan. For sake of illustration, if a user determines that a particular directory is used to store database event logs, and that such logs are not to be classified as sensitive, the user can define through an exclusion-based sampling criterion that the directory is to be skipped during scanning.

[0164] Another example of a sampling criterion includes a criterion based on one or more target risk signatures to be applied to identify instances of target risk events in the data store.

[0165] Further, as illustrated at block **1634**, the sampling criterion can include a combination of two or more of the criteria illustrated at block **1620-1632**. Of course, other sampling criterion can be utilized as well, as represented at block **1636**.

[0166] Alternatively, or in addition to user input to select the data store and/or sampling criterion, the data store and/or sampling criterion can be programmatically selected or otherwise identified, as represented at block **1638**. For instance, the sampling criterion can be identified based on the data store being scanned, as represented at block **1640**. For example, based on a volume and/or type of data in the data store, a sampling criterion can be programmatically selected. Alternatively, or in addition, the sampling criterion can be selected based on a previous scan of the data store, as represented as block **1642**. For instance, if a random sampling criterion was used in a previous scan of the data store, the same random sampling criterion can be used in one or more subsequent scans. In another example, a set of different sampling criterion can used rotatably, such that each subsequent scan of the data store uses a different sampling criterion than one or more prior scans of the store. Of course, the programmatic identification of the data store and/or sampling criterion can be performed in other ways as well, as represented at block **1644**.

[0167] At block **1646**, a data scanner is configured based on the selected sampling criterion. The data scanner can be configured by training a classifier (block **1648**) using training data that defines target data profiles (e.g., examples of sensitive data). Of course, the data scanner can be configured in other ways as well, as represented at block **1650**. At block **1652**, the data scanner is deployed on each data store.

[0168] The data scanner operates to scan the data store, which is represented by block **1654**. In instances in which multiple data stores are scanned by respective data scanners, the operations within block **1654** can be performed on each data store in parallel.

[0169] At block **1656**, a representative subset of data objects is selected from the data store based on the sampling criterion. As noted above, the sampling criterion can be selected based on manual input (block **1658**), automatically selected (block **1660**), selected based on a combination of manual and/or automatic processes (block **1662**), or selected in other ways as represented at block **1664**.

[0170] At block **1666**, instances of one or more pre-defined data patterns of the target data type are detected in the data store. For instance, block **1666** can identify a number of instances of sensitive and/or private data in the data store.

[0171] At block **1668**, a classifier result that collectively classifies the data store based on the sampling of the representative subset is generated. In one example, the classifier on the number of detected instances of the pre-defined data patterns to a threshold, as represented at block **1670**. In one example, the threshold can be set such that any number of detected instances of the target data type result in classification of the data store as having some correspondence to the target data type. For instance, a single instance of sensitive data in the data store can be utilized to classify the data store as having sensitive data. In another example, the threshold can be based on a percentage of the total number of data items in the data store. For example, the data store can be classified as sensitive if more than ten percent of the data items in the data store are identified as including sensitive data. This, of course, is for sake of example only.

[0172] At block **1672**, the classifier results can be output and, at block **1674**, one or more computing actions can be performed. In one example, at block **1676**, a user interface display can be generated to render an indication of the classifier result. For instance, the user interface display can identify the data stores identified as having sensitive data can be generated.

[0173] At block **1678**, risk signatures can be applied based on the classifier result. At block **1680**, one or more remedial actions can be performed. In one example, an infrastructure graph can be generated or updated as discussed above. Of course, other actions can be performed, as represented at block **1682**.

[0174] It can thus be seen that the present disclosure provides technology for data security posture analysis of a cloud environment. In some described examples, the technology implements intelligent data scanning sampling strategies, using sampling criterion, to identify a subset of data objects as being representative of a data store. A scanner analyzes only the representative subset of data object such that the data store is classified without having to scan the entire data set. This reduces the computational expense of data store classification while generating accurate classifications. Further, the data sampling is scalable as the data sets grow in size.

[0175] One or more implementations of the technology disclosed or elements thereof can be implemented in the form of a computer product, including a non-transitory computer readable storage medium with computer usable program code for performing the method steps indicated. Furthermore, one or more implementations and clauses of the technology disclosed or elements thereof can be implemented in the form of an apparatus including a memory and at least one processor that is coupled to the memory and operative to perform exemplary method steps. Yet further, in another aspect, one or more implementations and clauses of the technology disclosed or elements thereof can be implemented in the form of means for carrying out one or more of the method steps described herein; the means can include (i) hardware module(s), (ii) software module(s) executing on one or more hardware processors, or (iii) a combination of hardware and software modules; any of (i)-(iii) implement the specific techniques set forth herein, and the software modules are stored in a computer readable storage medium (or multiple such media).

[0176] Examples discussed herein include processor(s) and/or server(s). For sake of illustration, but not by limitation, the processors and/or servers include computer processors with associated memory and timing circuitry, and are functional parts of the corresponding systems or devices, and facilitate the functionality of the other components or items in those systems.

[0177] Also, user interface displays have been discussed. Examples of user interface displays can

take a wide variety of forms with different user actuatable input mechanisms. For instance, a user input mechanism can include icons, links, menus, text boxes, check boxes, etc., and can be actuated in a wide variety of different ways. Examples of input devices for actuating the input mechanisms include, but are not limited to, hardware devices (e.g., point and click devices, hardware buttons, switches, a joystick or keyboard, thumb switches or thumb pads, etc.) and virtual devices (e.g., virtual keyboards or other virtual actuators). For instance, a user actuatable input mechanism can be actuated using a touch gesture on a touch sensitive screen. In another example, a user actuatable input mechanism can be actuated using a speech command.

[0178] The present figures show a number of blocks with corresponding functionality described herein. It is noted that fewer blocks can be used, such that functionality is performed by fewer components. Also, more blocks can be used with the functionality distributed among more components. Further, the data stores discussed herein can be broken into multiple data stores. All of the data stores can be local to the systems accessing the data stores, all of the data stores can be remote, or some data stores can be local while others can be remote.

[0179] As used herein, if a description includes “one or more of” or “at least one of” followed by a list of example features with a conjunction “or” between the penultimate example feature and the last example feature, then this is to be read such that (1) one example includes at least one of or one or more of each feature of the listed features, (2) another example includes at least one of or one or more of only one feature of the listed features, and (3) another example includes some combination of the listed features that is less than all of the features and more than one of the features.

[0180] As used herein, if a description includes “one or more of” or “at least one of” followed by a list of example features with a conjunction “and” between the penultimate example feature and the last example feature, then this is to be read such that the example includes at least one of or one or more of each feature of all the listed features.

[0181] As used herein, if a description includes “one or more of” or “at least one of” followed by a list of example features with a conjunction “and/or” between the penultimate example feature and the last example feature, then this is to be read such that, in one example, the description includes “one or more of” or “at least one of” followed by a list of example features with a conjunction “or” between the penultimate example feature and the last example feature, and, in another example, the description includes “one or more of” or “at least one of” followed by a list of example features with a conjunction “and” between the penultimate example feature and the last example feature.

[0182] The above discussion has described a variety of different systems, components, logic, and interactions. One or more of these systems, components, logic and/or interactions can be implemented by hardware, such as processors, memory, or other processing components. Some particular examples include, but are not limited to, artificial intelligence components, such as neural networks, that perform the functions associated with those systems, components, logic, and/or interactions. In addition, the systems, components, logic and/or interactions can be implemented by software that is loaded into a memory and is executed by a processor, server, or other computing component, as described below. The systems, components, logic and/or interactions can also be implemented by different combinations of hardware, software, firmware, etc., some examples of which are described below. These are some examples of different structures that can be used to implement any or all of the systems, components, logic, and/or interactions described above.

[0183] The elements of the described figures, or portions of the elements, can be disposed on a wide variety of different devices. Some of those devices include servers, desktop computers, laptop computers, tablet computers, or other mobile devices, such as palm top computers, cell phones, smart phones, multimedia players, personal digital assistants, etc.

[0184] FIG. 17 is a simplified block diagram of one example of a client device **1700**, such as a handheld or mobile device, in which the present system (or parts of the present system) can be deployed. FIG. 18 illustrates an example of a handheld or mobile device.

[0185] One or more communication links **1702** allows device **1700** to communicate with other computing devices, and can provide a channel for receiving information automatically, such as by scanning. An example includes communication protocols, such as wireless services used to provide cellular access to a network, as well as protocols that provide local wireless connections to networks.

[0186] Applications or other data can be received on an external (e.g., removable) storage device or memory that is connected to an interface **1704**. Interface **1704** and communication links **1702** communicate with one or more processors **1706** (which can include processors or servers described with respect to the figures) along a communication bus (not shown in FIG. **17**), that can also be connected to memory **1708** and input/output (I/O) components **1710**, as well as clock **1712** and a location system **1714**.

[0187] Components **1710** facilitate input and output operations for device **1700**, and can include input components such as microphones, touch screens, buttons, touch sensors, optical sensors, proximity sensors, orientation sensors, accelerometers. Components **1710** can include output components such as a display device, a speaker, and or a printer port.

[0188] Clock **1712** includes, in one example, a real time clock component that outputs a time and date, and can provide timing functions for processor **1706**. Location system **1714** outputs a current geographic location of device **1700** and can include a global positioning system (GPS) receiver, a LORAN system, a dead reckoning system, a cellular triangulation system, or other positioning system. Memory **1708** stores an operating system **1716**, network applications and corresponding configuration settings **1718**, communication configuration settings **1720**, communication drivers **1722**, and can include other items **1724**. Examples of memory **1708** include types of tangible volatile and non-volatile computer-readable memory devices. Memory **1708** can also include computer storage media that stores computer readable instructions that, when executed by processor **1706**, cause the processor to perform computer-implemented steps or functions according to the instructions. Processor **1706** can be activated by other components to facilitate functionality of those components as well.

[0189] FIG. **18** illustrates one example of a tablet computer **1800** having a display screen **1802**, such as a touch screen or a stylus or pen-enabled interface. Screen **1802** can also provide a virtual keyboard and/or can be attached to a keyboard or other user input device through a mechanism, such as a wired or wireless link. Alternatively, or in addition, computer **1800** can receive voice inputs.

[0190] FIG. **19** shows an example computer system **1900** that can be used to implement the technology disclosed. Computer system **1900** includes at least one central processing unit (CPU) **1972** that communicates with a number of peripheral devices via bus subsystem **1955**. These peripheral devices can include a storage subsystem **1910** including, for example, memory devices and a file storage subsystem **1936**, user interface input devices **1938**, user interface output devices **1976**, and a network interface subsystem **1974**. The input and output devices allow user interaction with computer system **1900**. Network interface subsystem **1974** provides an interface to outside networks, including an interface to corresponding interface devices in other computer systems.

[0191] In one implementation, cloud security posture analysis system **1918** is communicably linked to the storage subsystem **1910** and the user interface input devices **1938**.

[0192] User interface input devices **1938** can include a keyboard; pointing devices such as a mouse, trackball, touchpad, or graphics tablet; a scanner; a touch screen incorporated into the display; audio input devices such as voice recognition systems and microphones; and other types of input devices. In general, use of the term “input device” is intended to include all possible types of devices and ways to input information into computer system **1900**.

[0193] User interface output devices **1976** can include a display subsystem, a printer, a fax machine, or non-visual displays such as audio output devices. The display subsystem can include an LED display, a cathode ray tube (CRT), a flat-panel device such as a liquid crystal display

(LCD), a projection device, or some other mechanism for creating a visible image. The display subsystem can also provide a non-visual display such as audio output devices. In general, use of the term “output device” is intended to include all possible types of devices and ways to output information from computer system **1900** to the user or to another machine or computer system.

[0194] Storage subsystem **1910** stores programming and data constructs that provide the functionality of some or all of the modules and methods described herein. These software modules are generally executed by processors **1978**.

[0195] Processors **1978** can be graphics processing units (GPUs), field-programmable gate arrays (FPGAs), application-specific integrated circuits (ASICs), and/or coarse-grained reconfigurable architectures (CGRAs). Processors **1978** can be hosted by a deep learning cloud platform such as Google Cloud Platform™, Xilinx™, and Cirrascale™. Examples of processors **1978** include Google's Tensor Processing Unit (TPU)™, rackmount solutions like GX4 Rackmount Series™, GX50 Rackmount Series™, NVIDIA DGX-1™, Microsoft's Stratix V FPGA™, Graphcore™ Intelligent Processor Unit (IPU)™, Qualcomm's Zeroth Platform™ with Snapdragon Processors™, NVIDIA's Volta™, NVIDIA's DRIVE PX™, NVIDIA's JETSON TX1/TX2 MODULE™, Intel's Nirvana™, Movidius VPU™, Fujitsu DPI™, ARM's DynamicIQ™, IBM TrueNorth™, Lambda GPU Server with Testa V100s™, and others.

[0196] Memory subsystem **1922** used in the storage subsystem **1910** can include a number of memories including a main random access memory (RAM) **1932** for storage of instructions and data during program execution and a read only memory (ROM) **1934** in which fixed instructions are stored. A file storage subsystem **1936** can provide persistent storage for program and data files, and can include a hard disk drive, a floppy disk drive along with associated removable media, a CD-ROM drive, an optical drive, or removable media cartridges. The modules implementing the functionality of certain implementations can be stored by file storage subsystem **1936** in the storage subsystem **1910**, or in other machines accessible by the processor.

[0197] Bus subsystem **1955** provides a mechanism for letting the various components and subsystems of computer system **1900** communicate with each other as intended. Although bus subsystem **1955** is shown schematically as a single bus, alternative implementations of the bus subsystem can use multiple busses.

[0198] Computer system **1900** itself can be of varying types including a personal computer, a portable computer, a workstation, a computer terminal, a network computer, a television, a mainframe, a server farm, a widely-distributed set of loosely networked computers, or any other data processing system or user device. Due to the ever-changing nature of computers and networks, the description of computer system **1900** depicted in FIG. **19** is intended only as a specific example for purposes of illustrating the preferred implementations of the present invention. Many other configurations of computer system **1900** are possible having more or less components than the computer system depicted in FIG. **19**.

[0199] It should also be noted that the different examples described herein can be combined in different ways. That is, parts of one or more examples can be combined with parts of one or more other examples. All of this is contemplated herein.

[0200] The technology disclosed can be practiced as a system, method, or article of manufacture. One or more features of an implementation can be combined with the base implementation. Implementations that are not mutually exclusive are taught to be combinable.

[0201] One or more features of an implementation can be combined with other implementations. This disclosure periodically reminds the user of these options. Omission from some implementations of recitations that repeat these options should not be taken as limiting the combinations taught in the preceding sections—these recitations are hereby incorporated forward by reference into each of the following implementations.

[0202] Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended

claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as example forms of implementing the claims.

Claims

1. A computer-implemented method for classifying a data store in a computing environment, the computer-implemented method comprising: identifying a first sampling criterion used for classification of the data store with respect to a target data type during a previous scan of the data store, wherein the data store stores a set of data objects; selecting a second sampling criterion, from a plurality of sampling criteria, based on the first sampling criterion used during the previous scan; deploying one or more scanners configured to select a subset of data objects, from the set of data objects stored in the data store, based on the second sampling criterion, wherein the subset comprises some, but not all, of the set of data objects; generating a classification result based on a number of instances of the target data type in the subset of data objects; and performing a computing action based on the classification result.
2. The computer-implemented method of claim 1, wherein the second sampling criterion is different than the first sampling criterion.
3. The computer-implemented method of claim 2, wherein the second sampling criterion comprises at least one of: a random sampling of the set of data objects from the data store; a time stamp-based sampling of the set of data objects based on time stamps representing a last modified time for each data object; a directory-based sampling of the set of data objects that is based on a directory structure in the data store; a metadata-based sampling of the set of data objects that is based on metadata of the set of data objects, the metadata comprising one or more of object type, object tags, or object size; or an exclusion-based sampling criterion.
4. The computer-implemented method of claim 2, wherein the first sampling criterion selected a different subset of the data objects in the data store than the subset of data objects selected based on the second sampling criterion.
5. The computer-implemented method of claim 1, wherein generating the classification result comprises classifying the data store as having a threshold correspondence to the target data type.
6. The computer-implemented method of claim 5, wherein classifying the data store comprises determining that the subset of data objects includes a threshold number of instances of one or more pre-defined data patterns representing the target data type.
7. The computer-implemented method of claim 1, wherein the computing environment comprises a cloud environment.
8. A computing system comprising: at least one processor; and memory storing instructions executable by the at least one processor, wherein the instructions, when executed, cause the computing system to: select a first sampling criterion representing a first distribution of data objects in a data store; select a second sampling criterion representing a second distribution of data objects in the data store; deploy one or more scanners configured to select a set of the data objects, in the data store that satisfy both the first sampling criterion and the second sampling criterion; obtain a scanner result, generated by the one or more scanners, that represents detected instances, in the set of the data objects, of one or more data patterns representing a target data type; generate a classification result based on detected instances of the one or more data patterns; and perform a computing action based on the classification result.
9. The computing system of claim 8, wherein each sampling criterion, of the first sampling criterion and the second sampling criterion, comprises at least one of: a random sampling criterion; a temporal sampling criterion based on last modified time stamps of the data objects; a directory-based sampling criterion based on a directory structure in the data store; a metadata-based sampling criterion that is based on metadata of the set of data objects; or an exclusion-based sampling criterion.

- 10.** The computing system of claim 8, wherein the first sampling criterion and the second sampling criterion are selected based on user input.
 - 11.** The computing system of claim 8, wherein the classification result identifies the data store as having a threshold correspondence to the target data type.
 - 12.** The computing system of claim 11, wherein the instructions, when executed, cause the computing system to determine that the set of data objects includes a threshold number of instances of the one or more data patterns.
 - 13.** The computing system of claim 11, wherein the one or more scanners are deployed in a cloud environment.
 - 14.** The computing system of claim 11, wherein first distribution of data objects is based on file type defined in metadata of the data objects.
 - 15.** A computer-readable media having computer-readable instructions stored thereon, wherein the computer-readable instructions, when executed by a computer, cause the computer to: identify a subject vulnerability signature representing an access path to target data in a computing environment; identify one or more characteristics associated with the target data; generate a sampling criterion based on the one or more characteristics of the target data; deploy one or more scanners configured to select a set of data objects in a data store based on the sampling criterion; generate a classification result based on an analysis of data objects, in the set of data objects, relative to the target data, the classification result representing a classification of the data store as having correspondence to the subject vulnerability signature; and perform a computing action based on the classification result.
 - 16.** The computer-readable media of claim 15, wherein the one or more characteristics comprise a target data type.
 - 17.** The computer-readable media of claim 16, wherein the sampling criterion represents a distribution of the set of data objects based on the target data type.
 - 18.** The computer-readable media of claim 15, wherein the classification result identifies the data store as having a threshold correspondence to the target data.
 - 19.** The computer-readable media of claim 18, wherein the computer-readable instructions, when executed by a computer, cause the computer to: determine that the set of data objects includes a threshold number of instances of the target data.
 - 20.** The computer-readable media of claim 15, wherein the computer-readable instructions, when executed by a computer, cause the computer to: receive a user input that selects the subject vulnerability signature.
-