

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250258851

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

Lewis; Daniel J. et al.

SYSTEMS AND METHODS FOR RETRIEVING TELEMATICS DATA

Abstract

The present disclosure relates to systems and methods for retrieving telematics data. A method for retrieving telematics data may comprise operating at least one processor to: provide at least one telematics database, the at least one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database; receive a natural language request from a user; generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request from the at least one telematics database; execute the executable query for retrieving the portion of the telematics data from the at least one telematics database; and return at least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database.

Inventors: Lewis; Daniel J. (Cambridge, CA), Mahdiani; Shadi (Brampton, CA), Zhang; Xin (Waterloo, CA), Poorsarvi Tehrani; Pooria (Holland Landing, CA), Weston; Kyle Micheal (Victoria, CA)

Applicant: Geotab Inc. (Oakville, CA)

Family ID: 94685734

Appl. No.: 19/052778

Filed: February 13, 2025

Related U.S. Application Data

us-provisional-application US 63553005 20240213

Publication Classification

Int. Cl.: G06F16/334 (20250101)

U.S. Cl.:

CPC G06F16/3344 (20190101); G06F16/3347 (20190101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application claims priority to and the benefit of U.S. Patent Application Ser. No. 63/553,005, filed on Feb. 13, 2024, which is hereby incorporated by reference in its entirety.

TECHNICAL FIELD

[0002] The present disclosure generally relates to the retrieval of telematics data from a database. More specifically, the present disclosure relates to retrieving telematics data based on a natural language request received from a user.

BACKGROUND

[0003] Today, many vehicles rely on computer-based systems (e.g., one or more processors) for their operation. As will be appreciated, such systems manage and/or produce many types of data associated with various aspects of the vehicle during the operation thereof that may generally be referred to as “telematics data”. In more detail, telematics data may include any information, parameters, attributes, characteristics, and/or features associated with the vehicle, such as, but not limited to, location data, speed data, acceleration data, fluid level data, energy data, engine data, brake data, transmission data, odometer data, vehicle identifying data, error/diagnostic data, tire pressure data, seatbelt data, and airbag data. The telematics data may be collected from the vehicle using, for example, a telematics device.

[0004] Telematics data may therefore include data relating to many different aspects a vehicle's operation. As will be appreciated, such a variety of information may be challenging to manage by a user, and in particular a user managing a vehicle fleet. For example, conventionally, a user may have to retrieve extended, comprehensive reports of the above-mentioned types of telematics data in order to gain insights about their vehicle or vehicle fleet. Such reports may be time-consuming to review, may require a lot of data to acquire (e.g., download), and may require a lot of space to store (e.g., in terms of bytes of a data storage). As well, it may be challenging to identify certain outliers, trends, patterns, etc. from such reports, particularly if the reports include data for many vehicles (e.g., a vehicle fleet).

[0005] A need therefore exists for improved systems and methods for retrieving telematics data.

SUMMARY

[0006] In one aspect, the present disclosure relates to a system for retrieving telematics data, the system comprising: at least one data storage operable to store at least one telematics database, the at least one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database; and at least one processor in communication with the at least one data storage, the at least one processor operable to: receive a natural language request from a user, the natural language request comprising at least one textual question relating to the telematics data stored within the at least one telematics database; generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is

responsive to the natural language request from the at least one telematics database by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request; and inputting the portion of contextual information into the LLM; execute the executable query for retrieving the portion of the telematics data from the at least one telematics database; and return at least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database.

[0007] According to an embodiment, the at least one processor is operable to identify the contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by the LLM via a keyword search.

[0008] According to a further embodiment, the context database is a vector database and the contextual information stored therein is represented by a plurality of vectors.

[0009] According to a further embodiment, the at least one processor is operable to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by: generating a vector representation of the parsing of the natural language request by the LLM; identifying one or more vectors of the plurality of vectors of the vector database that are similar to the vector representation of the parsing of the natural language request by the LLM; and selecting the contextual information represented by the one or more vectors that are similar to the vector representation for input into the LLM.

[0010] According to a further embodiment, the at least one processor is operable to identify the one or more vectors that are similar to the vector representation based on the real distance between the vector representation and each of the plurality of vectors of the vector database, based on an angle between the vector representation and the plurality of vectors of the vector database, or a combination thereof.

[0011] According to a further embodiment, the at least one data storage is further operable to store at least one chat database, the at least one chat database storing each natural language request received from the user.

[0012] According to a further embodiment, the at least one processor is further operable to identify additional contextual information based on one or more previous natural language requests received from the user.

[0013] According to a further embodiment, the at least one processor is further operable to modify the executable query based on database identifying information, a type of the portion of the telematics data that is responsive to the natural language request, an identity of the user, or a combination thereof.

[0014] According to a further embodiment, the at least one processor is further operable to revert modifying of the executable query after the execution thereof.

[0015] According to a further embodiment, the at least one processor is operable to return the portion of the telematics data and the executable query to the user.

[0016] According to a further embodiment, the at least one processor is further operable to: determine whether the executing of the executable query was successful in retrieving the portion of the telematics data; and generate, if the executing of the executable query was unsuccessful, an error message comprising a textual description of why the portion of telematics data was not retrieved.

[0017] According to a further embodiment, the at least one processor is further operable to: generate a corrected executable query for retrieving the portion of the telematics data from the

database by further inputting into the LLM the error message; and execute the corrected executable query for retrieving the portion of the telematics data from the database.

[0018] According to a further embodiment, the at least one processor is further operable to: determine whether the executing of the corrected executable query was successful in retrieving the portion of the telematics data; and repeat, if the executing of the corrected executable query was unsuccessful, the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query.

[0019] According to a further embodiment, the at least one processor is operable to repeat the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query until the executing of the corrected executable query retrieves the portion of the telematics data from the at least one telematics database.

[0020] According to a further embodiment, the at least one processor is operable to: repeat the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query a predetermined number of times; and return, if the repeating is performed the predetermined number of times without successfully retrieving the portion of the telematics data, a final error message to the user, the final error message comprising a textual description of why the portion of telematics data was not retrieved.

[0021] According to a further embodiment, final error message further comprises a textual request for an updated natural language request from the user.

[0022] According to a further embodiment, the at least one processor is further operable to send the natural language request and the executable query that is responsive thereto to the at least one data storage for storage in the context database.

[0023] According to a further embodiment, the at least one processor is further operable to receive from the user an indication of whether the executable query was responsive to the natural language request.

[0024] According to a further embodiment, the LLM comprises a generative artificial intelligence model.

[0025] According to a further embodiment, the one or more features of the database comprise information relating to how the telematics data is stored in the database, a type of telematics data stored in the database, or a combination thereof.

[0026] In another aspect, the present disclosure relates to a method for retrieving telematics data, the method comprising operating at least one processor to: provide at least one telematics database, the at least one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database; receive a natural language request from a user, the natural language request comprising at least one textual question relating to the telematics data stored within the at least one telematics database; generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request from the at least one telematics database by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request; and inputting the portion of contextual information into the LLM; execute the executable query for retrieving the portion of the telematics data from the at least one telematics database; and return at least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database.

[0027] According to an embodiment, the method comprises operating the at least one processor to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by the LLM via a keyword search.

[0028] According to a further embodiment, the at least one context database is a vector database and the contextual information stored therein is represented by a plurality of vectors.

[0029] According to a further embodiment, the method further comprises operating the at least one processor to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by: generating a vector representation of the parsing of the natural language request by the LLM; identifying one or more vectors of the plurality of vectors of the vector database that are similar to the vector representation of the parsing of the natural language request by the LLM; and selecting the contextual information represented by the one or more vectors that are similar to the vector representation for input into the LLM.

[0030] According to a further embodiment, the method comprises operating the at least one processor to identify the one or more vectors that are similar to the vector representation based on the real distance between the vector representation and each of the plurality of vectors of the vector database, based on an angle between the vector representation and the plurality of vectors of the vector database, or a combination thereof.

[0031] According to a further embodiment, the method further comprises operating the at least one processor to provide at least one chat database, the at least one chat database storing each natural language request received from the user.

[0032] According to a further embodiment, the method further comprises operating the at least one processor to identify additional contextual information based on one or more previous natural language requests received from the user.

[0033] According to a further embodiment, the method further comprises operating the at least one processor to modify the executable query based on database identifying information, a type of the portion of the telematics data that is responsive to the natural language request, an identity of the user, or a combination thereof.

[0034] According to a further embodiment, the method further comprises operating the at least one processor to revert modifying of the executable query after the execution thereof.

[0035] According to a further embodiment, the method comprises operating the at least one processor to return the portion of the telematics data and the executable query to the user.

[0036] According to a further embodiment, the method further comprises operating the at least one processor to: determine whether the executing of the executable query was successful in retrieving the portion of the telematics data; and generate, if the executing of the executable query was unsuccessful, an error message comprising a textual description of why the portion of telematics data was not retrieved.

[0037] According to a further embodiment, the method further comprises operating the at least one processor to: generate a corrected executable query for retrieving the portion of the telematics data from the database by further inputting into the LLM the error message; and execute the corrected executable query for retrieving the portion of the telematics data from the database.

[0038] According to a further embodiment, the method further comprises operating the at least one processor to: determine whether the executing of the corrected executable query was successful in retrieving the portion of the telematics data; and repeat, if the executing of the corrected executable query was unsuccessful, the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query.

[0039] According to a further embodiment, the method comprises operating the at least one processor to repeat the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query until the executing of the corrected

executable query retrieves the portion of the telematics data from the at least one telematics database.

[0040] According to a further embodiment, the method comprises operating the at least one processor to: repeat the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query a predetermined number of times; and return, if the repeating is performed the predetermined number of times without successfully retrieving the portion of the telematics data, a final error message to the user, the final error message comprising a textual description of why the portion of telematics data was not retrieved.

[0041] According to a further embodiment, the final error message further comprises a textual request for an updated natural language request from the user.

[0042] According to a further embodiment, the method further comprises operating the at least one processor to send the natural language request and the executable query that is responsive thereto to the at least one data storage for storage in the context database.

[0043] According to a further embodiment, the method further comprises operating the at least one processor to receive from the user an indication of whether the executable query was responsive to the natural language request.

[0044] According to a further embodiment, the LLM comprises a generative artificial intelligence model.

[0045] According to a further embodiment, the one or more features of the database comprise information relating to how the telematics data is stored in the database, a type of telematics data stored in the database, or a combination thereof.

[0046] In another aspect, the present disclosure relates to a non-transitory computer readable medium having instructions stored thereon executable by at least one processor to implement a method described herein.

[0047] Other aspects and features of the systems and methods of the present disclosure will become apparent to those ordinarily skilled in the art upon review of the following description of specific embodiments.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0048] These and other features of the present disclosure will become more apparent in the following detailed description in which reference is made to the appended drawings. The appended drawings illustrate one or more embodiments of the present disclosure by way of example only and are not to be construed as limiting the scope of the present disclosure.

[0049] FIG. 1 is a block diagram of various components interacting with an example fleet management system, according to an embodiment of the present disclosure.

[0050] FIG. 2 is a block diagram of an example fleet management system interacting with an example telematics device and an example vehicle, according to an embodiment of the present disclosure.

[0051] FIG. 3 is a block diagram of an example computing device interacting with an example fleet management system, according to an embodiment of the present disclosure.

[0052] FIG. 4 is a flowchart of an example method for retrieving telematics data, according to an embodiment of the present disclosure.

[0053] FIG. 5 is a block diagram of an example encoder-decoder transformer model, according to an embodiment of the present disclosure.

[0054] FIG. 6 is a block diagram of an example decoder-only transformer model, according to an embodiment of the present disclosure.

DETAILED DESCRIPTION

[0055] Telematics data may include a wide variety of different types of information, parameters, attributes, characteristics, features, and the like relating to various aspects of a vehicle. However, such a wide variety of data may be difficult to manage, especially if the telematics data is collected from a plurality of vehicles (e.g., a vehicle fleet).

[0056] For example, a user may wish to retrieve telematics data collected from their vehicle fleet to analyse a property or metric thereof. Conventionally, a user may have to retrieve a comprehensive report that details the telematics data collected from their vehicle fleet in order to analyse properties thereof. As described herein, such reports may include a significant amount of information and, as a result, may be time-consuming to review and difficult to parse and/or process if a user is inexperienced. For example, it may be challenging for some users to identify specific outliers, trends, patterns, etc. from such reports, particularly if the reports include data for many vehicles (e.g., a vehicle fleet). As well, due to the size of such reports, they may require a significant amount of data to obtain (e.g., via downloading) and a significant amount of space to store (e.g., in terms of bytes of a data storage).

[0057] It is therefore an object of the present disclosure to provide advantageous systems and methods for retrieving telematics data.

[0058] For example, in some embodiments, the systems and methods of the present disclosure may avoid the shortcomings of conventional techniques described above by retrieving only telematics data specifically requested by a user. In more detail, in such embodiments, a user may request a certain type of telematics data or, for example, a particular statistical analysis thereof, and the systems and methods described herein may be operable to return only that which was requested. As will be appreciated, such embodiments may reduce the amount of time spent parsing telematics data collected from a vehicle fleet by a user, as only specific, requested telematics data (or analyses thereof) may be returned. As well, by retrieving only particular types of telematics data, the overall size of the information returned may also be reduced, thereby decreasing the amount of data to obtain (e.g., download), as well as the space to store (e.g., in terms of bytes), the telematics data.

[0059] Further, as will be described herein, the systems and methods of the present disclosure may also retrieve the telematics data based on a natural language request from a user. As will be appreciated, a natural language request is a request that contains or is structured using “natural” language (i.e., a human language such as English) rather than an “artificial” or “constructed” language such as a computer programming language. By retrieving telematics data based on natural language requests, a user may request specific types of telematics data (or analyses thereof) without needing to be familiar with, for example, computer coding languages, how to execute complex statistical analyses, etc.

[0060] In light of the above, the systems and method of the present disclosure may provide a user a simple, efficient (e.g., in terms of time, data, and storage) way to access and process the telematics data collected from their vehicle or vehicle fleet.

[0061] Additional advantages will be discussed below and will be readily apparent to those of ordinary skill in the art upon reading the present disclosure.

[0062] Reference will now be made in detail to example embodiments of the disclosure, wherein numerals refer to like components, examples of which are illustrated in the accompanying drawings that further show example embodiments, without limitation.

[0063] Referring now to FIG. 1, there is shown an example of a fleet management system **110** for managing a plurality of assets equipped with a plurality of telematics devices **130**. Each of the telematics devices **130** is capable of collecting various data from the vehicles **120** (i.e., telematics data) and sharing the telematics data with the fleet management system **110**. The fleet management system **110** may be remotely located from the telematics devices **130** and the vehicles **120**.

[0064] The vehicles **120** may include any type of vehicle. For example, the vehicles **120** may include motor vehicles such as cars, trucks (e.g., pickup trucks, heavy-duty trucks such as class-8 vehicles, etc.), motorcycles, industrial vehicles (e.g., buses), and the like. Each motor vehicle may

be a gas, diesel, electric, hybrid, and/or alternative fuel vehicle. Further, the vehicles **120** may include vehicles such as railed vehicles (e.g., trains, trams, and streetcars), watercraft (e.g., ships and recreational pleasure craft), aircraft (e.g., airplanes and helicopters), spacecraft, and the like. Each of the vehicles **120** may be equipped with one of the telematics devices **130**.

[0065] Further, it is noted that, while only three vehicles **120** having three telematics devices **130** are shown in the illustrated example, it will be appreciated that there may be any number of vehicles **120** and telematics devices **130**. For example, the fleet management system **110** may manage hundreds, thousands, or even millions of vehicles **120** and telematics devices **130**.

[0066] In some embodiments, the telematics devices **130** may be standalone devices that are removably installed in the vehicles **120** (e.g., aftermarket telematics devices). In other embodiments, the telematics devices **130** may be integrated components of the vehicles **120** (e.g., pre-installed by an OEM). As described herein, the telematics devices **130** may collect various telematics data and share the telematics data with the fleet management system **110**. The telematics data may include any information, parameters, attributes, characteristics, and/or features associated with the vehicles **120**. For example, the telematics data may include, but is not limited to, location data, speed data, acceleration data, fluid level data (e.g., oil, coolant, and washer fluid), energy data (e.g., battery and/or fuel level), engine data, brake data, transmission data, odometer data, vehicle identifying data, error/diagnostic data, tire pressure data, seatbelt data, airbag data, or a combination thereof. In some embodiments, the telematics data may include information relating to the telematics devices **130** and/or other devices associated with or connected to the telematics devices **130**. Regardless, it should be appreciated the telematics data is a form of electronic data that requires a computer (e.g., a processor such as those described herein) to transmit, receive, interpret, process, and/or store.

[0067] Once received, the fleet management system **110** may process the telematics data obtained from the telematics devices **130** to provide various analysis, predictions, reporting, etc. In some embodiments, the fleet management system **110** may process the telematics data to provide additional information about the vehicles **120**, such as, but not limited to, trip distances and times, idling times, harsh braking and driving, usage rates, fuel economy, and the like. Various data analytics may be implemented to process the telematics data. The telematics data may then be used to manage various aspects of the vehicles **120**, such as route planning, vehicle maintenance, driver compliance, asset utilization, fuel management, etc., which, in turn, may improve productivity, efficiency, safety, and/or sustainability of the vehicles **120**.

[0068] A plurality of computing devices **150** may provide access to the fleet management system **110** to a plurality of users **160**. The users **160** may use computing devices **150** to access or retrieve various telematics data collected and/or processed by the fleet management system **110** to manage and track the vehicles **120**. As will be appreciated, the computing devices **150** may be any suitable computing devices. For example, the computing devices **150** may be any type of computers such as, but not limited to, personal computers, portable computers, wearable computers, workstations, desktops, laptops, smartphones, tablets, smartwatches, personal digital assistants (PDAs), mobile devices, and the like. The computing devices **150** may be remotely located from the fleet management system **110**, telematic devices **130**, and vehicles **120**.

[0069] The fleet management system **110**, telematics devices **130**, and computing devices **150** may communicate through a network **140**. The network **140** may comprise a plurality of networks and may be wireless, wired, or a combination thereof. As will be appreciated, the network **140** may employ any suitable communication protocol and may use any suitable communication medium. For example, the network **140** may comprise Wi-Fi™ networks, Ethernet networks, Bluetooth™ networks, near-field communication (NFC) networks, radio networks, cellular networks, and/or satellite networks. The network **140** may be public, private, or a combination thereof. For example, the network **140** may comprise local area networks (LANs), wide area networks (WANs), the internet, or a combination thereof. Of course, as will also be appreciated, the network **140** may also

facilitate communication with other devices and/or systems that are not shown.

[0070] Further, the fleet management system **110** may be implemented using one or more computers. For example, the fleet management system **110** may be implemented using one or more computer servers. The servers may be distributed across a wide geographical area. In some embodiments, the fleet management system **110** may be implemented using a cloud computing platform, such as Google Cloud Platform™ and Amazon Web Services™. In other embodiments, the fleet management system **110** may be implemented using one or more dedicated computer servers. In a further embodiment, the fleet management system **110** may be implemented using a combination of a cloud computing platform and one or more dedicated computer servers.

[0071] Referring now to FIG. 2, there is illustrated the fleet management system **110** in communication with one of the telematics devices **130** that is installed in one of the vehicles **120**. As shown, the fleet management system **110** may include a processor **112**, a data storage **114**, and a communication interface **116**, each of which may communicate with each other. The processor **112**, the data storage **114**, and the communication interface **116** may be combined into fewer components, divided into additional subcomponents, or a combination thereof. The components and/or subcomponents may not necessarily be distributed in proximity to one another and may instead be distributed across a wide geographical area.

[0072] The processor **112** may control the operation of the fleet management system **110**. As will be appreciated, the processor **112** may be implemented using one or more suitable processing devices or systems. For example, the processor **112** may be implemented using central processing units (CPUs), graphics processing units (GPUs), field programmable gate arrays (FPGAs), application specific integrated circuits (ASICs), digital signal processors (DSPs), neural processing units (NPU), quantum processing units (QPU), microprocessors, controllers, and the like. The processor **112** may execute various instructions, programs, software, or a combination thereof stored on the data storage **114** to implement various methods described herein. For example, the processor **112** may process various telematics data collected by the fleet management system **110** from the telematics devices **130**.

[0073] Various data for the fleet management system **110** may be stored on the data storage **114**. The data storage **114** may be implemented using one or more suitable data storage devices or systems such as random-access memory (RAM), read only memory (ROM), flash memory, hard disk drives (HDDs), solid-state drives (SSDs), magnetic tape drives, optical disc drives, memory cards, and the like. The data storage **114** may include volatile memory, non-volatile memory, or a combination thereof. Further, the data storage **114** may comprise non-transitory computer readable media. The data storage **114** may store various instructions, programs, and/or software that are executable by the processor **112** to implement various methods described herein. The data storage **114** may store various telematics data collected from the telematics devices **130** and/or processed by the processor **112**.

[0074] The communication interface **116** may enable communication between the fleet management system **110** and other devices and/or systems, such as the telematics devices **130**. The communication interface **116** may be implemented using any suitable communications devices and/or systems. For example, the communication interface **116** may comprise one or more various physical connectors, ports, or terminals such as universal serial bus (USB), ethernet, Thunderbolt, Firewire, serial advanced technology attachment (SATA), peripheral component interconnect (PCI), high-definition multimedia interface (HDMI), DisplayPort, and the like. As another example, the communication interface **116** may comprise one or more wireless interface components to connect to wireless networks such as Wi-Fi™, Bluetooth™, NFC, cellular, satellite, and the like. The communication interface **116** may enable various inputs and outputs to be received at and sent from the fleet management system **110**. For example, the communication interface **116** may be used to telematics data from the telematics devices **130**.

[0075] The telematics devices **130** also may include a processor **134**, a data storage **134**, and a

communication interface **136**. The telematics devices **130** may also comprise a sensor **138**. Each of the components of the telematics devices **130** may communicate with each other and may be combined into fewer components or divided into additional subcomponents.

[0076] The processor **132** may control the operation of the telematics device **130**. The processor **132** may be implemented using any suitable processing devices or systems, such as those described above in relation to the processor **112** of the fleet management system **110**. The processor **132** may execute various instructions, programs, software, or a combination thereof stored on the data storage **134** to implement various methods described herein. For example, the processor **132** may process various telematics data obtained from vehicle components **122** and/or the sensor **138**.

[0077] The data storage **134** may store various data for the telematics device **130**. The data storage **134** may be any suitable data storage device or system, such as those described above in relation to the data storage **114** of the fleet management system **110**. The data storage **134** may store various instructions, programs, software, or a combination thereof executable by the processor **132** to implement various methods described herein. As well, the data storage **134** may store various telematics data obtained from the vehicle components **122** and/or the sensor **138**.

[0078] The communication interface **136** may enable communication between the telematics devices **130** and other devices or systems, such as the fleet management system **110** and the vehicle components **122**. The communication interface **136** may comprise any suitable communication devices or systems, such as those described above in relation to the communication interface **116** of the fleet management system **110**. The communication interface **136** may enable various inputs and outputs to be received at and sent from the telematics devices **130**. For example, the communication interface **136** may be used to collect telematics data such as vehicle data from the vehicle components **122** and/or sensor **138**, to send telematics data to the fleet management system **110**, etc.

[0079] The sensor **138** may detect and/or measure various environmental events, changes, etc. The sensor **138** may include any suitable sensing devices or systems, such as, but not limited to, location sensors, velocity sensors, acceleration sensors, orientation sensors, vibration sensors, proximity sensors, temperature sensors, humidity sensors, pressure sensors, optical sensors, audio sensors, and combinations thereof. When the telematics device **130** is installed in the vehicle **120**, the sensor **138** may be used to collect telematics data that may not be obtainable from the vehicle components **122**. For example, the sensor **138** may include a satellite navigation device such as a global positioning system (GPS) receiver that may measure the location of the vehicle **120**. In some embodiments, the sensor **138** may comprise accelerometers, gyroscopes, magnetometers, inertial measurement units (IMUs), or the like that may measure the acceleration and/or orientation of the vehicle **120**.

[0080] In some embodiments, the telematics devices **130** may operate in conjunction with one or more accessory devices **170** that are in communication therewith. The accessory devices **170** may include one or more expansion devices that may provide additional functionality to the telematics devices **130**. For example, the accessory devices **170** may provide additional processing storage, communication, and/or sensing functionality through one or more additional processors, data storages, communication interfaces, and/or sensors (not pictured). The accessory devices **170** may also include adaptor devices that facilitate communication between the communication interface **136** and one or more vehicle interfaces **124**, such as a cable harness. The one or more accessory devices **170** may be installed in the vehicle **120** along with the telematics devices **130**.

[0081] As described herein, the telematics device **130** may be installed within the vehicle **120** removably or integrally. The vehicle **120** may include the vehicle components **122** and the one or more vehicle interfaces **124**, which, as will be appreciated, may be combined into fewer components or divided into additional subcomponents. In some embodiments, the vehicle components **122** may comprise any subsystems, parts, subcomponents, or combinations thereof of the vehicle **120**. For example, the vehicle components **122** may comprise powertrains, engines,

transmissions, steering, braking, seating, batteries, doors, suspensions, etc. The telematics device **130** may obtain various telematics data from the vehicle components **122**. For example, in some embodiments, the telematics device **130** may communicate with one or more electrical control units (ECUs) that control the vehicle components **122** or one or more internal sensors thereof.

[0082] The vehicle interface **124** may facilitate communication between the vehicle components **122** and other devices or systems. As well, the vehicle interface **124** may comprise any suitable communication devices or systems. For example, the vehicle interface **124** may include an on-board diagnostics (OBD-II) port and/or controller area network (CAN) bus port. The vehicle interface **124** may be used by the telematics device **130** to obtain telematics data from the vehicle components **122**. For example, the communication interface **136** may be connected to the vehicle interface **124** to communicate with the vehicle components **122**. In some embodiments, the one or more accessory devices **170** (e.g., a wire harness) may provide the connection between the communication interface **136** and the vehicle interface **124**.

[0083] Referring now to FIG. 3, there is shown the fleet management system **110** in communication with the computing devices **150**. As shown, the computing device **150** may also include a processor **152**, a data storage **153**, and a communication interface **156**. As well, the computing device **150** may include a display **158**. Each of the components of the computing device **150** may be communicate with each other and may be combined into fewer components or divided into additional subcomponents.

[0084] The processor **152** may control the operation of the computing device **150**. The processor **152** may be implemented using any suitable processing devices or systems, such as those described above in relation to the processor **112** of the fleet management system **110**. The processor **152** may execute various instructions, programs, software, or a combination thereof stored on the data storage **154** to implement various methods described herein. For example, the processor **152** may process various telematics data received from the fleet management system **110**, the telematics devices **130**, or a combination thereof.

[0085] The data storage **154** may store various data for the computing device **150**. The data storage **150** may be any suitable data storage device or system, such as those described above in relation to the data storage **114** of the fleet management system **110**. The data storage **154** may store various instructions, programs, software, or a combination thereof executable by the processor **152** to implement various methods described herein. As well, the data storage **154** may store various telematics data received from the fleet management system **110**, the telematics devices **130**, or a combination thereof.

[0086] The communication interface **156** may enable communication between the computing device **150** and other devices or systems, such as the fleet management system **110**. The communication interface **156** may be any suitable communication device or system, such as those described above in relation to the communication interface **116** of the fleet management system **110**. The communication interface **156** may enable various inputs and outputs to be received at and sent from the computing device **150**. For example, the communication interface **156** may be used to retrieve telematics data from the fleet management system **110**.

[0087] The displays **158** may visually present various data for the computing device **150**. The displays **158** may be implemented using any suitable display devices or systems, such as, but not limited to, light-emitting diode (LED) displays, liquid crystal displays (LCD), electroluminescent displays (ELDs), plasma displays, quantum dot displays, cathode ray tube (CRT) displays, and the like. The display **158** may be an integrated component that is integral with the computing device **150** or a standalone device that is removable connected to the computing device **150**. The display **158** may display various visual representations of the telematics data.

[0088] Referring now to FIG. 4, there is shown an example of a method for retrieving telematics data (**400**) according to an embodiment of the present disclosure. As shown, the method **400** comprises operating at least one processor to: provide at least one telematics database, the at least

one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database (**410**); receive a natural language request from a user, the natural language request comprising at least one textual question relating to the telematics data stored within the at least one telematics database (**420**); generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request from the at least one telematics database by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request; and inputting the portion of contextual information into the LLM (**430**); execute the executable query for retrieving the portion of the telematics data from the at least one telematics database (**440**); and return at least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database (**450**).

[0089] The method **400** may be implemented using any suitable combination of hardware and software, such as those described in reference to FIG. **1** to FIG. **3**. For example, one or more operations (e.g., operations **410**, **420**, **430**, **440**, and/or **450**) of the method **400** may be implemented at the fleet management system (e.g., by the processor **112** executing instructions stored on the data storage **114**), at the telematics device **130** (e.g., by the processor **132** executing instructions stored on the data storage **134**), at the computing devices **150** (e.g., by the processor **152** executing instructions stored on the data storage **154**), or a combination thereof.

[0090] At operation **410** of the method **400**, at least one telematics database and at least one context database may be provided. The at least one telematics database may store telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and the at least one context database may store contextual information relating to the at least one telematics database.

[0091] Each of the at least one telematics database and the at least one context database may be any type of database that is suitable for storing and accessing the data stored therein. For example, each of the at least one telematics database and the at least one context database may be relational databases, which, as will be appreciated, store information in tables, rows, and columns and may be queried using common computer coding languages such as structured query language (SQL). Each of the at least one telematics database and the at least one context database may independently be located on a data storage (e.g., the data storage **114**, **134**, and/or **154**), may be cloud-based (e.g., Cloud SQL), or a combination thereof. Each of the at least one telematics database and the at least one context database may be located on the same data storage, or different data storages if so desired.

[0092] The data stored in each of the telematics database and the contextual database may be organized therewithin using any suitable system. For example, each of the at least one telematics database may include telematics data originating from telematics devices installed in the vehicles of a single user. As another example, each context database may include contextual information relating to a particular type of telematics data, a particular telematics database, etc.

[0093] With respect to the at least one telematics database, the telematics data stored therein may be obtained from a plurality of vehicles using, for example, the systems outlined in FIG. **1** to FIG. **3**. For instance, the telematics data may originate from the telematics devices **130** installed in the plurality of vehicles **120**. As described herein, the telematics data may generally include information, parameters, attributes, characteristics, and/or features associated with the vehicle. The

telematics data may be “raw” (i.e., unprocessed) and/or processed data. That is, the telematics data may be obtained from the telematics devices installed in the vehicles and input directly into the at least one telematics database. Additionally, or alternatively, the telematics data may be processed (e.g., by the fleet management system **110**) to provide additional information about the vehicles from which it originates and then included in the at least one telematics database.

[0094] In addition to those described above, examples of telematics data (e.g., raw and/or processed telematics data) include, but are not limited to, location data, speed data, acceleration data, fluid level data, energy data, engine data, brake data, transmission data, odometer data, vehicle identifying data, error/diagnostic data, tire pressure data, seatbelt data, airbag data, vehicle identification numbers, telematics device identifying information, telematics device health information, odometer data (e.g., the highest odometer value measured within the timeframe), operational times (e.g., the total time that the vehicle was actively in operation), idle times (e.g., the total time the vehicle was idling), ignition times (e.g., the total time that the vehicle engine was in ignition), distances travelled, number of stops made, total fuel consumed (e.g., during a trip and/or while idling), fuel economy estimates, fault code information (e.g., the number and/or types of vehicle and/or engine fault codes measured), telematics device fault code information, the latest measured latitudinal and longitudinal values, trip distances and times, idling times, harsh braking and driving, usage rates, fuel economy, safety-related predictions, predicted maintenance events (e.g., time-to-service), etc.

[0095] By a variety of types of telematics data within the at least one telematics database, the telematics data may be more readily retrieved (e.g., via an executable query) and returned to a user, as the databases may include less information (e.g., uncommonly requested types of telematics data and/or data analytics), thereby reducing the total size of the thereof, and a processor (e.g., one or more of the processors **112**, **132**, **152**) may not have to process the data upon receipt of a request from the user (e.g., the curated telematics data may include pre-processed telematics data as described above). As will be appreciated, the overall efficiency of the systems and methods of the present disclosure may, as a result, be increased.

[0096] As indicated herein, the at least one context database may store contextual information relating to the telematics data stored within the at least one telematics database. Such contextual information may include, for example, one or more features of the at least one telematics database (e.g., how the telematics data is stored, the types of telematics data stored, data analytics functions stored in the at least one telematics database, etc.). As well, and as will be discussed below, the contextual information may include a plurality of example natural language requests and corresponding executable query outputs that may exemplify to a machine learning model how to generate an executable query for accessing the at least one telematics database. Additional types of contextual information that may be stored in the at least one context database will be discussed herein.

[0097] Referring back to FIG. **4**, at operation **420**, a natural language request from a user may be received. The natural language request may comprise at least one textual question relating to the telematics data stored within one of the plurality of databases. That is, the textual question may be structured in “natural language”, which, as described herein, refers to ordinary human language (e.g., “plain” English), rather than an “artificial” or “constructed” language such as a computer programming language.

[0098] The textual question may relate to the telematics data stored on the database in that the question is answerable thereby. For example, the textual question may be about a certain type of telematics data, one or more properties of a vehicle or vehicle fleet informed by the telematics data, data analytics performed on the telematics data, etc. Examples of natural language textual questions that relate to telematics data include, but are not limited to, “How much did my fleet idle last month?”, “How does the fuel economy of our older vehicles (based on the year of manufacture) compare to our newer vehicles?”, “What is the fleet's average fuel economy over the past 6

months?”, “Assuming a gas price of \$3.5 per gallon, how much money have we lost due to idling in the last 6 months?”, “Which vehicles have the lowest idle to drive ratio in the past month?”, “Which group is an outlier in terms of idling over the past 6 months?”, “How have EVs been adopted in my fleet over the last year”, etc.

[0099] Of course, it will be appreciated that the textual question may not always be formatted by a user as a question. For example, the textual question may instead be structured as an answerable command. For example, the textual question may be a command such as “List my top 10 vehicles with the lowest utilization over the past week”, “List the top 10 groups from highest to lowest idling in the past week”, “Provide the vehicles having the lowest idle to drive ratio in the past month”, etc. Thus, the textual question may be a question or command that is answerable by the telematics data stored in the at least one telematics database.

[0100] The natural language request may be received from the user via any suitable system. For example, a user may input the textual question of the natural language request into a form, a field, or any other suitable interface using a computing device (e.g., the computing devices **150**). The textual question may then be sent to a processor (e.g., one or more of the processors **112**, **132**, **152**) via a communication interface (e.g., the communication interface **156**) for further processing thereby. In some embodiments, the methods and systems of the present disclosure may be implemented as a chatbot (i.e., may mimic having a conversation with a user during use thereof). In such embodiments, the user may input the natural language request as a portion of a conversation with the chatbot.

[0101] As shown at operation **430**, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request may be generated. As used herein, an “executable query” generally refers to computer code that, when executed, is capable of retrieving information from a database. Thus, in the context of the systems and methods of the present disclosure, the executable query, when executed, may retrieve the portion of the telematics data from the database. The executable query may be generated in any suitable computer programming language compatible with the plurality of databases. For example, the executable query may be generated in SQL. In some embodiments, the executable query may be a wrapped executable query to facilitate the execution thereof immediately after generation.

[0102] Further, in some embodiments, the executable query may include portions that are replaceable with information specific to the database. In more detail, as described above, the machine learning model will generally not have access to the plurality of databases. Thus, it may in some cases be necessary to insert into the executable query after the generation thereof, certain database identifying information so that the executable query, when executed, may access and retrieve the information stored in the databases. Examples of such database identifying information include, database names, user IDs associated with particular databases, etc.

[0103] The portion of the telematics data that is to be retrieved by the executable query may be responsive to the natural language request of the user in that the portion of the telematics data “answers” the textual question of the natural language request. For example, if the user requests a specific type of telematics data obtained from their vehicle fleet (i.e., a portion of the total telematics data included in the databases), the executable query may retrieve that type of telematics data from the relevant database.

[0104] Thus, to respond to the natural language request, the executable query may be generated such that it at least retrieves a portion of the telematics data. In some embodiments, depending on the natural language request, the executable query may also process the retrieved telematics data. For example, the executable query may be generated to retrieve and perform one or more mathematical operations to the telematics data such as, but not limited to, one or more statistical analyses. In such embodiments, the processed retrieved portion of telematics data may be responsive to the natural language request.

[0105] The executable query may be generated using a machine learning model. The machine

learning model may be a model that is capable of converting natural language into computer programming language. The machine learning model may be, for example, an artificial neural network such as a transformer machine learning model. As will be appreciated, such transformer models are often used for processing (or “transforming”) natural language prompts. Examples of transformer models include encoder-only models, encoder-decoder models, and decoder-only models. In the example embodiment of FIG. 4, an LLM is used to generate the executable query. As will be appreciated, LLMs are typically large-scale implementations of such transformer models.

[0106] For illustrative purposes, FIG. 5 shows a simplified block diagram of an example encoder-decoder transformer model **500**. As shown, the transformer model **500** may include: an input layer **510**, an encoder component **520**, a decoder component **530**, a linear layer **540**, and a softmax layer **550**.

[0107] The input layer **510** generally prepares an input for processing by the transformer model **500**. The input may be a textual input (e.g., the natural language request) and may sometimes be referred to as a “prompt”. The input layer **510** may comprise an input embeddings sublayer **512** and a positional encoding sublayer **514**. The input embeddings sublayer **512** may receive the textual input, tokenize the textual input to generate tokens that each correspond to a portion of the textual input (e.g., a word or a portion of a word of a sentence), and generate a vector embedding of the tokens as vectors, the vectors indicating the semantic meaning and/or contextual information of the tokens. The positional encoding sublayer **514** may then positionally encode the vector embedding to include the relative position of each of the tokens within the input (e.g., where words of a sentence are positioned relative to each other and/or within the sentence).

[0108] The encoder component **520** generally generates encodings that indicate which tokens are relevant to one another. In more detail, the encoder component **520** may comprise a plurality of encoding layers **522** (for simplicity, only one encoding layer **522** is shown), each of which may comprise a multi-head self-attention sublayer **524** and a feed-forward neural network layer **526**. The multi-head self-attention sublayer **524** may receive the vector embedding having the positional information included therein and assign a weight to each of the tokens represented thereby based on their relevance to other tokens (e.g., via a scaled dot-product attention function), thereby generating an attention embedding. Using multi-head attentions, each token may be assigned a weight multiple times in parallel, which may then be averaged (e.g., a weighted average) to determine the assigned weight for each thereof. After the multi-head attention sublayer **524**, the attention embedding may be received by the feed-forward neural network sublayer **526**, which may perform a plurality of linear regressions to transform the attention embedding into an encoded output for further processing by the transformation model **500** (e.g., the another encoding layer **522** and/or the decoder component **530**).

[0109] The decoder component **530** generally predicts tokens of an output sequentially (i.e., one-at-a-time) based at least in part on encoded output of the encoder component **530**. The decoder component **530** may comprise a plurality of decoding layers **532** (for simplicity, only one decoding layer **532** is shown), each of which may comprise a masked multi-head self-attention sublayer **534**, a multi-head self-attention sublayer **536**, and a feed-forward neural network layer **538**. The masked multi-head self-attention layer **534** and the multi-head self-attention sublayer **536** function substantially the same as the multi-head self-attention sublayer **524** of the encoder component **520**. However, it is noted that the masked multi-head self-attention layer **534** also receives each sequential output of the transformer model **500** (e.g., each predicted token as they are predicted) and masks any subsequent tokens such that only earlier (i.e., with respect to an output sentence) tokens are considered when assigning weights, thereby generating a masked output. The multi-head self-attention sublayer **536** may receive both the encoded output from the encoder component **520** as well as the masked output from the masked multi-head self-attention layer **534** to generate a decoded attention embedding. After the masked multi-head self-attention layer **534** and the multi-

head self-attention sublayer **536**, the decoded attention embedding may be received by the feed-forward neural network layer **538**, which may function substantially the same as the feed-forward neural network sublayer **526**, to produce a decoder output.

[0110] The decoder output may then be received by linear layer **540**, which performs a linear transformation to generate a logits vector based on the decoder output. The softmax layer **550** then converts the logits vector into predicted next-token probabilities, and the next-token associated with a highest probability is selected as the output of the transformation model **500**.

[0111] If the output of the transformation model **500** indicates that the response is complete, the process may end. Otherwise, the output of the transformation model **500** may be input into the decoder component **530** (i.e., at a masked multi-head self-attention layer **532**) and the process may continue. As will be appreciated, if the output of the transformation model **500** is in a natural language, it may first be input into a second input layer **560**, comprising an output embeddings sublayer **562** and a positional encoding sublayer **564**, which may function substantially the same as the input embeddings sublayer **512** and the positional encoding sublayer **514**, respectively, prior to the decoder component **530**.

[0112] For further illustration, another transformer model **600** is illustrated in FIG. **6**. The transformer model **600** is a decoder-only model and is structured similarly to the decoder component **530** of the transformer model **500**. In more detail, the transformer model **600** may include an input layer **610**, a decoder component **620**, a linear layer **630**, and a softmax layer **640**.

[0113] The input layer **610** may comprise an input embeddings sublayer **612** and a positional encoding sublayer **614**, which function substantially the same as the input embeddings sublayer **512** and the positional encoding sublayer **514**, respectively, of the transformer model **500** to generate a vector embedding having the positional information of the tokens included therein.

[0114] The decoder component **620** may comprise a plurality of decoding layers **622** (for simplicity, only one decoding layer **622** is shown), each of which may comprise a masked multi-head self-attention sublayer **624** and a feed-forward neural network layer **626**. The masked multi-head self-attention sublayer **624** and the feed-forward neural network layer **626** function substantially the same as the masked multi-head self-attention sublayer **536** and the feed-forward neural network layer **538** of the transformer model **500**, respectively. Thus, the decoding layer **622** may not include a non-masked multi-head self-attention sublayer (e.g., the multi-head self-attention sublayers **524**, **536**). In more detail, the vector embedding output by the input layer **610** is received by the masked multi-head self-attention sublayer **624** and processed through the feed-forward neural network layer **626** and any subsequent decoding layers **622**. Each output is generated iteratively (i.e., one-at-a-time) and with reference only to previous outputs.

[0115] The linear layer **630** and the softmax layer **640** also function substantially the same as the linear layer **540** and the softmax layer **550**, respectively, of the transformer model **500**. That is, the linear layer **630** and the softmax layer **640** process the output from the decoder component **620** to select a next-token having a highest probability as the output of the transformer model **600**.

[0116] Similar to the transformer model **500**, if the output of the transformation model **600** indicates that the response to the input is complete, the process may end. Otherwise, the output of the transformation model **600** may be input into the decoder component **620** (i.e., at a masked multi-head self-attention layer **624**) and the process may continue. As will be appreciated, if the output of the transformation model **600** is in a natural language, it may first be re-input into the input layer **610** prior to the decoder component **620**.

[0117] It is noted that transformer models such as the transformer models **500**, **600** may be particularly well-suited for parallel implementations, allowing for faster outputs, reduced training times, etc. Due to these characteristics, transformer models are particularly scalable and are capable of being trained on large datasets of text (e.g., text scrapped various internet websites, databases, etc.). As will be appreciated, transformer models are therefore often used as the underlying architectures for LLMs.

[0118] Referring back to operation 430 of FIG. 4, the particular machine learning model employed to generate the executable query may be implemented using any suitable system such as one or more of the systems described in relation to FIG. 1 to FIG. 3 (e.g., operated via one or more of the processors 112, 132, 152).

[0119] Further, as indicated above, machine learning models such as LLMs may be pre-trained using large amounts of public data (e.g., text scraped from internet websites) so that they may predict an appropriate output based on an input. However, such models generally do not have access to data contained in private databases—e.g., databases that store a user's confidential telematics data. It may therefore be challenging to use a machine learning model to generate an executable query that is executable to retrieve data (e.g., telematics data) from a private database (i.e., a database that is not accessible to the machine learning model). However, the inventors of the present disclosure surprisingly found that, by providing sufficient context to a machine learning model such as an LLM, it may be capable of generating an executable query that, when executed, may retrieve data from a database inaccessible thereto.

[0120] Thus, the executable query for retrieving the telematics data may be generated using, for example, an LLM that does not have access to the at least one telematics database storing thereon the telematics data, by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request; and inputting the portion of contextual information into the LLM.

[0121] As indicated above in relation to FIG. 5 and FIG. 6, when an input is input into an LLM, the machine learning model may “parse” the input to identify relationships between the components thereof—e.g., grammatical relationships, syntactical relationships, and/or the like between portions of the natural language request. Thus, the machine learning model (e.g., an LLM) may “parse” the natural language request when input therein to identify the meaning, or intent, thereof.

[0122] By parsing the natural language request, the machine learning model may determine information required to generate the executable query. As described above, the machine learning model generally may not have access to the at least one telematics database storing thereon the telematics data and, as a result, may not have access to the structure of the at least one telematics database, the types of telematics data stored in the at least one telematics database, etc. The machine learning model will therefore typically need contextual information so that it can generate an executable query capable of retrieving the portion of telematics data from the at least one telematics database that is responsive to the natural language request.

[0123] Thus, when generating the executable query, the at least one processor may be operated to identify, based on a parsing of the natural language request by the machine learning model (e.g., an LLM) a portion of the contextual information useful for generating the executable query from the at least one context database, as described above. In more detail, the portion of contextual information may comprise one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request.

[0124] The one or more features of the at least one telematics database may include information such as, but not limited to, how the telematics data is stored in the database, the types of telematics data stored in the database, data analytics functions stored in the at least one telematics database, etc. For example, in some embodiments, the one or more features of the at least one telematics database may comprise table identifying information (e.g., table names indicating the telematics data organized therein) and the schema thereof. In more detail, the schema may include information relating to the telematics data included in the table, the format of that telematics data (e.g., string,

timestamp, float, etc.), and a description of the telematics data, structured in a way that is accessible via the executable query.

[0125] The one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request may comprise, for instance, examples of input natural language requests and corresponding generated executable queries that were, or will be, successful in retrieving the appropriate portion of telematics data (i.e., the telematics data that was responsive to the natural language requests) from the at least one telematics database. The one or more example natural language requests and corresponding executable query outputs may include natural language requests previously input by a user and the corresponding previously generated executable queries, as will be described herein. Alternatively, or additionally, the one or more example natural language requests and corresponding executable query outputs may comprise those generated by, for example, an administrator of the systems and methods of the present disclosure that would be successful in retrieving the appropriate portion of telematics data.

[0126] As will be appreciated, the one or more example natural language requests and corresponding executable query outputs may include those relating to a variety of types of telematics data. For example, the one or more example natural language requests and the corresponding executable query outputs may include those relating to vehicle idling, vehicle distances travelled, vehicle fuel consumption, and/or other miscellaneous queries. Further, as indicated above, the one or more one or more example natural language requests and corresponding executable query outputs may be relevant to the natural language request. That is, the one or more example natural language requests and corresponding executable query outputs may, for example, relate to the same type of telematics data, require a similar processing step (e.g., normalization of telematics data, aggregation of telematics data obtained within a particular time period, and the like), etc.

[0127] The one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request may be identified using a number of techniques. For example, in some embodiments, the at least one processor may be operated to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by the machine learning model (e.g., an LLM) via a keyword search. In such embodiments, the at least one processor may identify the one or more example natural language requests and corresponding executable query outputs based on, for example, the presence of a same, or similar, word or words in the natural language request and an example natural language request. As will be appreciated, in such embodiments, the at least one context database may be operable to store the contextual information as plain, searchable text.

[0128] In some embodiments, the at least one context database may be a vector database, and the contextual information stored therein may be represented by a plurality of vectors. As will be appreciated, and as indicated above in relation to FIG. 5 and FIG. 6, a vector database may store information as vectors, wherein the position, magnitude, and orientation of each vector may represent the content (e.g., textual content) and/or “meaning” (e.g., based on grammatical, syntactical, and/or textual relationships) of a particular portion of the information. Thus, in embodiments where the at least one context database is a vector database, using a keyword search to identify one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request may not be particularly useful. In such embodiments, it may instead be useful to operate the at least one processor to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by, for example: generating a vector representation of the parsing of the natural language request by the machine learning model (e.g., a LLM); identifying one or more vectors of the plurality of vectors of the vector database that are similar to the vector representation of the parsing of the natural language request by the machine

learning model; and selecting the contextual information represented by the one or more vectors that are similar to the vector representation for input into the machine learning model.

[0129] The vector representation of the parsing of the natural language request may be generated using any suitable technique (e.g., a pre-trained sentence embedding algorithm). As described herein, the parsing of the natural language request by the machine learning model may indicate textual, grammatical, and/or syntactical relationships thereof. As also described herein, such relationships may affect the positioning, the magnitude, the orientation, etc. of a vector corresponding to the parsing in the vector database.

[0130] Once the vector representation of the parsing of the natural language request has been generated, one or more vectors of the plurality of vectors of the vector database that are similar to the vector representation may be identified. The vector representation of the parsing may be determined to be similar to one or more vectors of the plurality of vectors using any suitable technique. For example, in some embodiments, the at least one processor (e.g., one or more of the processors **112**, **132**, **152**) may be operated to identify the one or more vectors that are similar to the vector representation based on the real distance (e.g., Euclidian distance) between the vector representation and each of the plurality of vectors of the vector database, based on an angle between the vector representation and the plurality of vectors of the vector database, or a combination thereof. For instance, in such embodiments, a vector of the vector database that is closest to the vector representation of the parsing in terms of real distance and/or angle of orientation may be selected as a vector that is similar to the vector representation, and the portion of contextual information represented thereby may be selected for input into the machine learning model. Of course, it may in some cases be useful to select a plurality of vectors as similar to the vector representation so that the identified portion of contextual information provided to the machine learning model contains as much useful information as possible. In such cases, the at least one processor may be operated to, for example, select the two, three, four, or so on, most similar of the vectors of the vector database (e.g., based on real distance and/or angle therebetween).

[0131] As indicated above, once the portion of the contextual information useful for generating the executable query from the at least one context database is identified, it may be input into the machine learning model so that the executable query may be generated thereby.

[0132] Further, as indicated above, in some embodiments, the systems and methods of the present disclosure may be implemented as a “chat bot”—i.e., the systems and methods of the present disclosure may receive and return information to a user in the style of a natural language conversation. In such embodiments, multiple, consecutive natural language requests may be received from a user. Due to how natural language conversations typically transpire (e.g., between humans), it may be the case that one or more of the consecutive natural language requests received from a user may require context provided by a previous natural language request and/or a previously returned portion of telematics data to generate an executable query for returning a portion of the telematics data that is responsive thereto. In such cases, it may be useful to store (permanently or temporarily) any previously received natural language requests and corresponding executable query outputs for use as additional contextual information by the machine learning model.

[0133] For example, in some embodiments, the method **400** may further comprise operating the at least one processor to provide at least one chat database, the chat database storing each natural language request received from the user. In some embodiments, the chat database may store each response to the natural language request by the systems and methods of the present disclosure (e.g., the returned portion of telematics data). The at least one processor may then identify additional contextual information based on one or more previous natural language requests received from the user, for input into the machine learning model (e.g., an LLM) for the generation of the executable query.

[0134] The inventors surprisingly found that, by inputting a portion of the contextual information

useful for generating the executable query (i.e., that includes one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request), the machine learning model may more reliably generate an executable query capable of retrieving telematics data that is responsive to the natural language request without needing access to the at least one telematics database. Thus, the systems and methods of the present disclosure may be particularly well-suited for mitigating security and/or privacy related issues, as databases containing private information (e.g., the at least one telematics database) are not directly accessed by the machine learning model. [0135] At operation **440** of FIG. **4**, the executable query for retrieving the portion of the telematics data from the at least one telematics database may be executed. The executable query may be executed using any suitable combinations of hardware and software such as, for example, those described herein in relation to FIG. **1** to FIG. **3**. Of course, the particular execution of the executable query may depend at least in part on, for example, the structure of the executable query (e.g., the computer programming language that the executable query is generated in).

[0136] In some embodiments, the method **400** may further comprise operating the at least one processor to modify the executable query after the generation thereof. For example, the executable query may be modified based on database identifying information, a type of the portion of the telematics data that is responsive to the natural language request, an identity of the user, or a combination thereof. For example, the executable query may be modified to include an identity of the user such that only the database or databases storing telematics data originating from that user's telematics devices are accessible by the executable query. As another example, if the type of the portion of telematics data that is responsive to the natural language request is a type that is often spelled incorrectly, or has different names in different regions (e.g., gas, gasoline, petrol, fuel, etc.), the executable query may be modified to account for such discrepancies.

[0137] However, in some embodiments, the executable query may be returned to the user. In such cases, it may be useful to remove any modifications to the executable query so as to omit user identifying information (e.g., the ID of the user), database identifying information that may be confidential for the administrator implementing the systems and methods described herein, or any other information included in the executable query during the modification thereof that might be confidential or unnecessary for a user to view. Thus, the method **400** may further comprise operating the at least one processor to revert the modifying of the executable query after the execution thereof. The executable query may be reverted, for example, to its original state—i.e., as generated by the machine learning model.

[0138] In some cases, it may be useful to determine whether the executable query successfully retrieved the portion of telematics data that is responsive to the natural language request. For example, in some embodiments, the method **400** may further comprise operating the at least one processor to determine whether the executing of the executable query was successful in retrieving the position of the telematics data; and generate, if the executing of the executable query was unsuccessful, an error message comprising at least a textual description of why the telematics data was not retrieved.

[0139] The determining of whether the executable query was successful in retrieving the portion of the telematics data may be implemented using any suitable system or technique. For example, the at least one processor (e.g., one or more of the processors **112**, **132**, **152**) may be operable to check one or more of whether the plurality of data bases were accessed, whether the correct database storing the portion of telematics data was accessed, whether any telematics data was retrieved, or any other indication that the portion of telematics data was retrieved, or not retrieved. Of course, the exact technique for determining whether the telematics data was retrieved may vary based at least in part on the particular implementation of the system.

[0140] If the at least one processor is unable to determine any indication that the portion of telematics data was retrieved, the error message may be generated. As described above, the error

message may comprise at least a textual description of why the telematics data was not retrieved. For example, the error message may explain that the database storing the portion of telematics data was not accessed, that the telematics data was not located at the location included in the executable query, or any other reason why the telematics data was not retrieved. As will be appreciated, an exact reason may differ system to system, based at least in part on the implementation thereof. [0141] Once the error message is generated, it may be returned to the user and/or used to adjust the executable query. For example, in some embodiments, the method **400** may further comprise operating the at least one processor to: generate a corrected executable query for retrieving the portion of the telematics data from the database by inputting into the machine learning model (e.g., an LLM) the error message; and execute the corrected executable query for retrieving the portion of the telematics data from the database. In such embodiments, by inputting the error message, additional context may be provided to the machine learning model. For example, the machine learning model, may use the error message to adjust or “correct” the executable query based on the textual description of why the telematics data was not retrieved included therein so as to generate a corrected executable query that is capable of retrieving the portion of the telematics data that is responsive to the natural language request from the at least one telematics database.

[0142] Once the corrected executable query is generated, it may be executed to retrieve the telematics data. However, it may in some cases be useful to again check if the corrected executable query retrieves the portion of the telematics data that is responsive to the natural language request. For example, in some embodiments, the method **400** may further comprise operating the at least one processor to determine whether the executing of the corrected executable query was successful in retrieving the portion of the telematics data. The determining of whether the execution of the corrected executable query was successful may be implemented using any suitable system or technique, such as those described above in relation to determining whether the executable query was successful.

[0143] If the corrected executable query is determined to be unsuccessful, the above process may be repeated, or not. For example, in some embodiments, if the executing of the corrected executable query was unsuccessful, the method **400** may further comprise operating the at least one processor to repeat the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query. In such embodiments, the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query until the executing of the corrected executable query retrieves the portion of the telematics data from the database—e.g., until the corrected executable query is determined to have been successful.

[0144] In another embodiment, if the executing of the corrected executable query was unsuccessful, the method **400** may further comprise operating the at least one processor to repeat the generating of the error message, the generating of the corrected executable query, and the executing of the corrected executable query a predetermined number of times; and return, if the repeating is performed the predetermined number of times without successfully retrieving the portion of the telematics data, a final error message to the user, the final error message comprising at least a textual description of why the telematics data was not retrieved. That is, in such embodiments, the operations to generate a successful corrected query may be limited to being repeated the predetermined number of times. Such embodiments may be useful to reduce processing times so that a response may be more quickly returned to a user. The predetermined number of times that the repeating of the operations for generating the corrected executable query may be any suitable number and is not particularly limited. For example, the operations for generating the corrected executable query may be repeated 5, 10, 15, or 20 times or any number more or fewer or therebetween if so desired.

[0145] As described above, if a successful corrected executable query is not generated after repeating the operations for the generation thereof the predetermined number of times, a final error

message comprising a textual description of why the telematics data was not retrieved may be returned to the user. In some embodiments, the final error message may further comprise a textual request for an updated natural language request from the user. In such embodiments, the user may be prompted to input an updated natural language request comprising, for example, a modified or different textual question relating to the telematics data stored on the database so that the methods and systems of the present disclosure may re-attempt to retrieve the desired telematics data.

[0146] Referring back to FIG. 4, as shown at operation 450, once retrieved (e.g., by the executable query or the corrected executable query), at least the portion of the telematics data that is responsive to the natural language request may be returned to the user. The portion of the telematics data may be returned using any suitable system, such as those described above in relation to FIG. 1 to FIG. 3. For example, the processor 152 of the computing device 150 may return the telematics data to the user such that it is viewable on the display 158 thereof. As well, the telematics data may be returned in any suitable format. For example, the telematics data may be returned in a table, as a single line of information, as a natural language response, etc.

[0147] In some embodiments, additional information or content may be returned to the user with the portion of telematics data. For example, the executable query (or the corrected executable query, as the case may be) may be returned to the user to provide additional information to the user about how the telematics data was retrieved. In such embodiments, it may be desirable to revert any modifications made to the executable query prior to the returning thereof to the user, as previously described herein. As another example, additionally or alternatively, a natural language response may be returned with the telematics data, the natural language response comprising, for example, a textual description of the returned telematics data, a textual response to the textual question of the natural language request that comprises the telematics data, etc. Such configurations may be useful if implementing the systems and methods of the present disclosure as a chatbot, as described above.

[0148] It is also noted that the portion of the telematics data that is responsive to the natural language request may be returned to the user without providing the machine learning model (e.g., an LLM) access to the telematics data stored on the plurality of databases, as indicated above. As will be appreciated, machine learning models may retain information processed thereby to inform future outputs. While useful for providing more accurate outputs over time, if confidential information is entered as an input, the machine learning model may permanently retain that confidential information. The systems and methods of the present disclosure avoid such vulnerabilities, as the machine learning model is used to generate the executable query but not access a user's confidential telematics data. As a result, the systems and methods of the present disclosure may be implemented using third-party machine learning models such as publicly available LLMs (e.g., generative artificial intelligence models) without risk of confidential data being retained thereby.

[0149] Further, as described above, the portion of contextual information input into the machine learning model (e.g., an LLM) may include one or more example natural language requests and corresponding executable query outputs. In some embodiments, the one or more example natural language requests and corresponding executable query outputs may be pre-generated or pre-selected (e.g., by an administrator of the systems and methods of the present disclosure). However, additionally or alternatively, in some embodiments, the one or more example natural language requests may be those that were previously received by one or more users and the corresponding executable query outputs may be those previously generated in response to the previously received natural language requests. In such embodiments, the method 400 may further comprise operating the at least one processor to receive from the user an indication of whether the executable query was responsive to the natural language request. If the user indicates that the executable query was responsive to the natural language request, the method 400 may further comprise operating the at least one processor to include the natural language request and the executable query in the contextual information stored in the at least one context database.

[0150] Further, as a machine learning model (e.g., an LLM) may be used to generate the executable query, it may be desirable to train, or pre-train, the machine learning model using examples of natural language requests and corresponding executable query outputs. As will be appreciated, a machine learning model may be trained by inputting training data for processing to, over time, recognize patterns, relationships, etc. therein. As described above, in some embodiments, the method **400** may further comprise operating the at least one processor to send the natural language request and the executable query indicated by the user to have been successful in retrieving telematics data that was responsive to the natural language request to the at least one context database for storage therein. In such embodiments, the natural language request and the executable query may also be included in training data for training the machine learning model. As will be appreciated, by training the machine learning model, the output executable queries may be generated more quickly and more accurately. As well, by training the machine learning model, less information may need to be provided thereto by way of the portion of contextual information, which may increase processing efficiency and thus scalability of the systems and methods of the present disclosure.

[0151] In the present disclosure, all terms referred to in singular form are meant to encompass plural forms of the same. Likewise, all terms referred to in plural form are meant to encompass singular forms of the same. Unless defined otherwise, all technical and scientific terms used herein have the same meaning as commonly understood by one of ordinary skill in the art to which this disclosure pertains.

[0152] As used herein, the term “about” refers to an approximately $\pm 10\%$ variation from a given value. It is to be understood that such a variation is always included in any given value provided herein, whether or not it is specifically referred to.

[0153] It should be understood that the compositions and methods are described in terms of “comprising,” “containing,” or “including” various components or steps, the compositions and methods can also “consist essentially of or “consist of the various components and steps. Moreover, the indefinite articles “a” or “an,” as used in the claims, are defined herein to mean one or more than one of the element that it introduces.

[0154] Throughout this specification and the appended claims, infinitive verb forms are often used, such as “to operate” or “to couple”. Unless context dictates otherwise, such infinitive verb forms are used in an open and inclusive manner, such as “to at least operate” or “to at least couple”.

[0155] For the sake of brevity, only certain ranges are explicitly disclosed herein. However, ranges from any lower limit may be combined with any upper limit to recite a range not explicitly recited, as well as, ranges from any lower limit may be combined with any other lower limit to recite a range not explicitly recited, in the same way, ranges from any upper limit may be combined with any other upper limit to recite a range not explicitly recited. Additionally, whenever a numerical range with a lower limit and an upper limit is disclosed, any number and any included range falling within the range are specifically disclosed. In particular, every range of values (of the form, “from about a to about b,” or, equivalently, “from approximately a to b,” or, equivalently, “from approximately a-b”) disclosed herein is to be understood to set forth every number and range encompassed within the broader range of values even if not explicitly recited. Thus, every point or individual value may serve as its own lower or upper limit combined with any other point or individual value or any other lower or upper limit, to recite a range not explicitly recited.

[0156] The Drawings are not necessarily to scale and may be illustrated by phantom lines, diagrammatic representations, and fragmentary views. In certain instances, details that are not necessary for an understanding of the exemplary embodiments or that render other details difficult to perceive may have been omitted.

[0157] The specification includes various implementations in the form of block diagrams, schematics, and flowcharts. A person of skill in the art will appreciate that any function or operation within such block diagrams, schematics, and flowcharts can be implemented by a wide

range of hardware, software, firmware, or combination thereof. As non-limiting examples, the various embodiments herein can be implemented in one or more of: application-specific integrated circuits (ASICs), standard integrated circuits (ICs), programmable logic devices (PLDs), field-programmable gate arrays (FPGAs), computer programs executed by any number of computers or processors, programs executed by one or more control units or processor units, firmware, or any combination thereof.

[0158] The disclosure includes descriptions of several processors. Said processors can be implemented as any hardware capable of processing data, such as application-specific integrated circuits (ASICs), standard integrated circuits (ICs), programmable logic devices (PLDs), field-programmable gate arrays (FPGAs), logic circuits, or any other appropriate hardware. The disclosure also includes descriptions of several non-transitory processor-readable storage mediums. Said non-transitory processor-readable storage mediums can be implemented as any hardware capable of storing data, such as magnetic drives, flash drives, RAM, or any other appropriate data storage hardware. Further, mention of data or information being stored at a device generally refers to the data information being stored at a non-transitory processor-readable storage medium of said device.

[0159] Therefore, the present disclosure is well adapted to attain the ends and advantages mentioned as well as those that are inherent therein. The particular embodiments disclosed above are illustrative only, as the present disclosure may be modified and practiced in different but equivalent manners apparent to those skilled in the art having the benefit of the teachings herein. Although individual embodiments are discussed, the disclosure covers all combinations of all those embodiments. Furthermore, no limitations are intended to the details of construction or design herein shown, other than as described in the claims below. Also, the terms in the claims have their plain, ordinary meaning unless otherwise explicitly and clearly defined by the patentee. It is therefore evident that the particular illustrative embodiments disclosed above may be altered or modified and all such variations are considered within the scope and spirit of the present disclosure. If there is any conflict in the usages of a word or term in this specification and one or more patent(s) or other documents that may be incorporated herein by reference, the definitions that are consistent with this specification should be adopted.

[0160] Many obvious variations of the embodiments set out herein will suggest themselves to those skilled in the art in light of the present disclosure. Such obvious variations are within the full intended scope of the appended claims.

Claims

1. A system for retrieving telematics data, the system comprising: at least one data storage operable to store at least one telematics database, the at least one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database; and at least one processor in communication with the at least one data storage, the at least one processor operable to: receive a natural language request from a user, the natural language request comprising at least one textual question relating to the telematics data stored within the at least one telematics database; generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request from the at least one telematics database by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that

are relevant to the natural language request; and inputting the portion of contextual information into the LLM; execute the executable query for retrieving the portion of the telematics data from the at least one telematics database; and return at least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database.

2. The system of claim 1, wherein the context database is a vector database and the contextual information stored therein is represented by a plurality of vectors.

3. The system of claim 2, wherein the at least one processor is operable to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by: generating a vector representation of the parsing of the natural language request by the LLM; identifying one or more vectors of the plurality of vectors of the vector database that are similar to the vector representation of the parsing of the natural language request by the LLM; and selecting the contextual information represented by the one or more vectors that are similar to the vector representation for input into the LLM.

4. The system of claim 3, wherein the at least one processor is operable to identify the one or more vectors that are similar to the vector representation based on the real distance between the vector representation and each of the plurality of vectors of the vector database, based on an angle between the vector representation and the plurality of vectors of the vector database, or a combination thereof.

5. The system of claim 1, wherein the at least one data storage is further operable to store at least one chat database, the at least one chat database storing each natural language request received from the user.

6. The system of claim 5, wherein the at least one processor is further operable to identify additional contextual information based on one or more previous natural language requests received from the user.

7. The system of claim 1, wherein the at least one processor is further operable to modify the executable query based on database identifying information, a type of the portion of the telematics data that is responsive to the natural language request, an identity of the user, or a combination thereof.

8. The system of claim 7, wherein the at least one processor is further operable to revert modifying of the executable query after the execution thereof.

9. The system of claim 1, wherein the at least one processor is further operable to send the natural language request and the executable query that is responsive thereto to the at least one data storage for storage in the context database.

10. A method for retrieving telematics data, the method comprising operating at least one processor to: provide at least one telematics database, the at least one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database; receive a natural language request from a user, the natural language request comprising at least one textual question relating to the telematics data stored within the at least one telematics database; generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request from the at least one telematics database by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request; and inputting the portion of contextual information into the LLM; execute the executable query for retrieving the portion of the telematics data from the at

least one telematics database; and return at least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database.

11. The method of claim 10, wherein the at least one context database is a vector database and the contextual information stored therein is represented by a plurality of vectors.

12. The method of claim 11, comprising operating the at least one processor to identify the portion of contextual information useful for generating the executable query from the at least one context database based on the parsing of the natural language request by: generating a vector representation of the parsing of the natural language request by the LLM; identifying one or more vectors of the plurality of vectors of the vector database that are similar to the vector representation of the parsing of the natural language request by the LLM; and selecting the contextual information represented by the one or more vectors that are similar to the vector representation for input into the LLM.

13. The method of claim 12, comprising operating the at least one processor to identify the one or more vectors that are similar to the vector representation based on the real distance between the vector representation and each of the plurality of vectors of the vector database, based on an angle between the vector representation and the plurality of vectors of the vector database, or a combination thereof.

14. The method of claim 10, further comprising operating the at least one processor to provide at least one chat database, the at least one chat database storing each natural language request received from the user.

15. The method of claim 14, further comprising operating the at least one processor to identify additional contextual information based on one or more previous natural language requests received from the user.

16. The method of claim 10, further comprising operating the at least one processor to modify the executable query based on database identifying information, a type of the portion of the telematics data that is responsive to the natural language request, an identity of the user, or a combination thereof.

17. The method of claim 16, further comprising operating the at least one processor to revert modifying of the executable query after the execution thereof.

18. The method of claim 10, further comprising operating the at least one processor to send the natural language request and the executable query that is responsive thereto to the at least one data storage for storage in the context database.

19. A non-transitory computer readable medium having instructions stored thereon executable by at least one processor to implement a method for retrieving telematics data, the method comprising operating at least one processor to: provide at least one telematics database, the at least one telematics database storing telematics data originating from a plurality telematics devices installed in a plurality of vehicles, and at least one context database, the at least one context database storing contextual information relating to the at least one telematics database; receive a natural language request from a user, the natural language request comprising at least one textual question relating to the telematics data stored within the at least one telematics database; generate, using a large language model (LLM) that does not have access to the at least one telematics database, an executable query for retrieving a portion of the telematics data that is responsive to the natural language request from the at least one telematics database by: inputting the natural language request into the LLM; identifying, based on a parsing of the natural language request by the LLM, a portion of the contextual information useful for generating the executable query from the at least one context database, the portion of the contextual information comprising one or more features of the at least one telematics database and one or more example natural language requests and corresponding executable query outputs that are relevant to the natural language request; and inputting the portion of contextual information into the LLM; execute the executable query for retrieving the portion of the telematics data from the at least one telematics database; and return at

least the portion of the telematics data to the user, whereby the natural language request is responded to without providing the LLM with access to the telematics data stored on the at least one telematics database and the at least one context database.
