US Patent & Trademark Office
Patent Public Search | Text View

United States Patent Application Publication   20250259384
Kind Code   A1
Publication Date   August 14, 2025
Inventor(s)   Wald; Ingo

MULTI-DIMENSIONAL BINNING FOR HIERARCHICAL PARTITIONING

Abstract

In various examples, spatial elements of a scene may be projected into multi-dimensional bins that correspond to a spatial partitioning of the scene to determine assignments between the spatial elements and the multi-dimensional bins. A partition of the spatial elements may be determined using the assignments and a spatial element may be assigned to a node corresponding to a hierarchical partitioning of the spatial elements based on the partition. To determine the partition, candidate split planes may be determined with respect to the multi-dimensional bins, and a split plane that defines the partition may be selected from the candidate split planes. The assignments and the multi-dimensional bins may also be used to determine subpartitions of the partition. For example, the assignments may be used to determine the subpartitions with respect to a subset of the multi-dimensional bins that corresponds to the partition.

Inventors:	**Wald; Ingo (Salt Lake City, UT)**
Applicant:	**NVIDIA Corporation** (Santa Clara, CA)
Family ID:	**96661246**
Appl. No.:	**18/437071**
Filed:	**February 08, 2024**

Publication Classification

Int. Cl.:   **G06T17/00** (20060101)

U.S. Cl.:

CPC   **G06T17/005** (20130101); G06T2210/21 (20130101)

Background/Summary

BACKGROUND

[0001] Bounding Volume Hierarchies (BVHs) may be used to accelerate ray intersection testing with scene (e.g., object) geometries for light transport simulation techniques when rendering images of scenes. BVHs are used to organize the geometry of a scene into a hierarchical tree-like structure, where a node may represent a bounding volume that encloses a portion of the scene. As a ray traverses a hierarchy of a BVH, for rays that do not intersect with a node's bounds, the subtree beneath that node (representing a region of the scene) in the hierarchy can be discarded from further testing, thereby reducing the number of intersection tests used to test for an intersection with the geometry of the scene. When the scene geometry changes, a new BVH may need to be built or a previous BVH may be modified to reflect updated scene conditions. However, a slow building process for a BVH can be a bottleneck in performing the light transport simulation.

[0002] A typical approach to building a BVH for a frame involves a recursive process in which a set of primitives for a node are partitioned into multiple subsets. The partitioning may be achieved using a split plane that separates the primitives into sets for new nodes, then repeating the process on those nodes until a leaf node is reached. Each time a split plane is to be determined, the primitives may be read from memory, and for each axis (e.g., an x-axis, a y-axis, and a z-axis), projected into bins or buckets along the axis to determine split plane candidates along that axis. The split plane for a node may then be selected from the split plane candidates from multiple axes using a cost function, such as a Surface Area Heuristic (SAH). However, such a process may be performed for each frame of a scene as the viewport of the viewer/user changes; as scenes often include a large quantity of primitives, this approach may consume significant memory bandwidth, which may greatly reduce the build speed for the BVH.

SUMMARY

[0003] Embodiments of the present disclosure relate to multi-dimensional binning for hierarchical partitioning. Systems and methods are disclosed that may be used to reduce the bandwidth required to read spatial elements from memory to determine hierarchical partitionings of spatial elements of a scene.

[0004] In contrast to conventional systems, such as those described above, spatial elements of a scene may be projected into multi-dimensional bins that correspond to a spatial partitioning of the scene to determine assignments between the spatial elements and the multi-dimensional bins. One or more partitions of the spatial elements may then be determined based at least on one or more values indicating the assignments, and at least one spatial element of the spatial elements may be assigned to a node corresponding to a hierarchical partitioning of the spatial elements based at least on the partition. In at least one embodiment, a partition may be determined based at least on selecting one or more split planes that define the partition. For example, one or more candidate split planes may be determined with respect to one or more dimensions of the multi-dimensional bins. The candidate split planes may be evaluated using the one or more values, and the one or more split planes may be selected based at least on the evaluation. In at least one embodiment, the assignments and the multi-dimensional bins may be used to determine one or more subpartitions of the one or more partitions. For example, the assignments may be used to determine the one or more subpartitions with respect to a subset of the multi-dimensional bins that corresponds to a partition. Similar to a partition, a subpartition may be determined based at least on selecting one or more split planes that define the subpartition(s).

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] The present systems and methods for multi-dimensional binning for hierarchical partitioning are described in detail below with reference to the attached drawing figures, wherein:

[0006] FIG. **1** is an example of a process for determining a hierarchical partitioning of spatial elements using multi-dimensional binning, in accordance with some embodiments of the present disclosure;

[0007] FIG. **2** is an example of a diagram illustrating bounding shapes which may be associated with a hierarchical partitioning of spatial elements, in accordance with some embodiments of the present disclosure;

[0008] FIG. **3** is an example of a diagram illustrating projection of spatial elements into multi-dimensional bins, in accordance with some embodiments of the present disclosure;

[0009] FIG. **4** is an example of a diagram illustrating split planes which may be determined using multi-dimensional bins, in accordance with some embodiments of the present disclosure;

[0010] FIG. **5** is a flow diagram showing a method for projecting spatial elements into multi-dimensional bins to select split planes for partitions of spatial elements, in accordance with embodiments of the present disclosure;

[0011] FIG. **6** is a flow diagram showing a method for determining values indicating assignments between spatial elements and multi-dimensional bins to select split planes for partitions of spatial elements, in accordance with embodiments of the present disclosure;

[0012] FIG. **7** illustrates an example parallel processing unit suitable for use in implementing at least some embodiments of the present disclosure;

[0013] FIG. **8**A illustrates an example general processing cluster within the parallel processing unit of FIG. **7** suitable for use in implementing at least some embodiments of the present disclosure;

[0014] FIG. **8**B illustrates an example memory partition unit of the parallel processing unit of FIG. **7** suitable for use in implementing at least some embodiments of the present disclosure;

[0015] FIG. **9**A illustrates an example of the streaming multi-processor of FIG. **8**A suitable for use in implementing at least some embodiments of the present disclosure;

[0016] FIG. **9**B is an example conceptual diagram of a processing system implemented using the PPU of FIG. **7** suitable for use in implementing at least some embodiments of the present disclosure;

[0017] FIG. **9**C illustrates an example system in which the various architecture and/or functionality of the various embodiments may be implemented;

[0018] FIG. **10** illustrates an example ray tracing pipeline suitable for use in implementing at least some embodiments of the present disclosure;

[0019] FIG. **11** illustrates an example acceleration structure suitable for use in implementing at least some embodiments of the present disclosure;

[0020] FIG. **12** illustrates an example shader record suitable for use in implementing at least some embodiments of the present disclosure;

[0021] FIG. **13** is a block diagram of an example computing device suitable for use in implementing some embodiments of the present disclosure; and

[0022] FIG. **14** is a block diagram of an example data center suitable for use in implementing some embodiments of the present disclosure.

DETAILED DESCRIPTION

[0023] Systems and methods are disclosed related to multi-dimensional binning for hierarchical partitioning. Disclosed approaches may be used to reduce the bandwidth required to read spatial elements from memory to determine hierarchical partitionings of spatial elements of a scene.

[0024] In accordance with disclosed approaches, spatial elements (e.g., primitives) of a scene may be projected into multi-dimensional bins that correspond to a spatial partitioning of the scene to determine assignments between the spatial elements and the multi-dimensional bins. One or more

partitions of the spatial elements may then be determined based at least on one or more values indicating the assignments and at least one spatial element of the spatial elements may be assigned to a node corresponding to a hierarchical partitioning of the spatial elements based at least on the partition. In at least one embodiment, a partition may be determined based at least on selecting one or more split planes that define the partition. For example, one or more candidate split planes may be determined with respect to one or more dimensions of the multi-dimensional bins (e.g., along an X-axis, a Y-axis, or a Z-axis). The candidate split planes may be evaluated using the one or more values, such as counters of the quantity of spatial elements assigned to the bins, merged bounding boxes of the spatial elements assigned to the bins, etc., and the one or more split planes may be selected based at least on the evaluation.

[0025] In at least one embodiment, the assignments and the multi-dimensional bins may be used to determine one or more subpartitions of the one or more partitions. For example, the one or more values indicating the assignments may be used to determine the one or more subpartitions with respect to a subset of the multi-dimensional bins that corresponds to partition. Similar to a partition, a subpartition may be determined based at least on selecting one or more split planes that define the subpartition(s). For example, one or more candidate split planes may be determined with respect to one or more dimensions of the subset of the multi-dimensional bins (e.g., along an X-axis, a Y-axis, or a Z-axis). The candidate split planes may be evaluated using the one or more values, such as counters of the quantity of spatial elements assigned to the bins, merged bounding boxes of the spatial elements assigned to the bins, etc., and the one or more split planes may be selected based at least on the evaluation.

[0026] In at least one embodiment, one or more spatial elements that correspond to a subpartition of a partition of spatial elements may be assigned to a child node to a parent node that corresponds to the partition in the hierarchical partitioning of the spatial elements. Additionally, or in the alternative, the one or more spatial elements that correspond to the subpartition of the partition of spatial elements may be assigned to a child node to a parent node that corresponds to the spatial elements.

[0027] The hierarchical partitioning may be used to perform any of a variety of light transport simulation operations on graphical data, such as path tracing, ray tracing, ray marching, etc. By way of example, and not limitation, the hierarchical partitioning may correspond to an acceleration structure, such as a Bounding Volume Hierarchy (BVH) using for intersection testing in light transport simulation. However, the hierarchical partitioning may be used for additional or alternative applications.

[0028] The systems and methods described herein may be used for a variety of purposes, by way of example and without limitation, for machine control, machine locomotion, machine driving, synthetic data generation, model training, perception, augmented reality, virtual reality, mixed reality, robotics, security and surveillance, autonomous or semi-autonomous machine applications, deep learning, environment simulation, object or actor simulation and/or digital twinning, data center processing, conversational AI, light transport simulation (e.g., ray tracing, path tracing, etc.), collaborative content creation for 3D assets, cloud computing, generative AI, (large) language models, and/or any other suitable applications.

[0029] Disclosed embodiments may be comprised in a variety of different systems such as automotive systems (e.g., a control system for an autonomous or semi-autonomous machine, a perception system for an autonomous or semi-autonomous machine), systems implemented using a robot, aerial systems, medial systems, boating systems, smart area monitoring systems, systems for performing deep learning operations, systems for performing simulation operations, systems for performing digital twin operations, systems implemented using an edge device, systems incorporating one or more virtual machines (VMs), systems for performing synthetic data generation operations, systems implemented at least partially in a data center, systems for performing conversational AI operations, systems for performing light transport simulation,

systems for performing collaborative content creation for 3D assets, systems implemented at least partially using cloud computing resources, systems for performing generative AI operations, systems for performing operations using a large language model, and/or other types of systems.

[0030] With reference to FIG. **1**, FIG. **1** is an example of a process **100** for determining a hierarchical partitioning of spatial elements using multi-dimensional binning, in accordance with some embodiments of the present disclosure. It should be understood that this and other arrangements described herein are set forth only as examples. Other arrangements and elements (e.g., machines, interfaces, functions, orders, groupings of functions, etc.) may be used in addition to or instead of those shown, and some elements may be omitted altogether. Further, many of the elements described herein are functional entities that may be implemented as discrete or distributed components or in conjunction with other components, and in any suitable combination and location. Various functions described herein as being performed by entities may be carried out by hardware, firmware, and/or software. For instance, various functions may be carried out by a processor executing instructions stored in memory. In at least one embodiment, the systems, methods, and processes described herein may be executed using similar components, features, and/or functionality to those of the parallel processing unit **700** of FIG. **7**, the general processing clusters **750** of FIG. **8**A, the memory partition unit **780** of FIG. **8**B, the streaming multiprocessor **840** of FIG. **9**A, the processing system **900** of FIG. **9**B, the system **965** of FIG. **9**C, the computing device **1300** of FIG. **13**, and/or the data center **1400** of FIG. **14**.

[0031] The process **100** may be implemented using, amongst additional or alternative components, one or more node managers **102**, one or more bin determiners **104**, one or more bin assignors **106**, or one or more partition determiners **108**.

[0032] As an overview, the node manager **102** may be configured to manage nodes of a hierarchical partitioning of a scene, such as a hierarchical partitioning **120**, which may include modifying, creating, and/or deleting one or more nodes. The bin determiner **104** may be configured to determine and/or generate multi-dimensional bins that correspond to a spatial partitioning of the scene, such as multi-dimensional bins **122**. The bin assignor **106** may be configured to assign one or more spatial elements to one or more of the multi-dimensional bins, such as a spatial element **124**A, a spatial element **124**B, a spatial element **124**C, or a spatial element **124**D (also referred to as "spatial elements **124**"), for example, based at least on projecting the spatial elements into the multi-dimensional bins. The partition determiner **108** may be configured to determine one or more partitions of the spatial elements based at least on one or more assignments between the one or more spatial elements and the one or more multi-dimensional bins determined using the bin assignor **106**. The node manager **102** may be configured to assign at least one of the spatial elements to a node of the hierarchical partitioning or otherwise update or generate the hierarchical partitioning based at least on the one or more partitions determined using the partition determiner **108**. In at least one embodiment, the process **100** is implemented using a recursive process for determining the hierarchical partitioning **120**. The recursive process may include recursively partitioning spatial elements of a node to determine child nodes for the hierarchical partitioning **120** until a leaf node is determined and/or an end condition it satisfied for the recursion.

[0033] A hierarchical partitioning determined using the process **100** may be used to render an image of a scene using any of a variety of rendering techniques. For example, the hierarchical partitioning **120** may be used to perform any of a variety of light transport simulation operations on graphical data, such as path tracing, ray tracing, ray marching, etc. By way of example, and not limitation, the hierarchical partitioning may correspond to the acceleration structure **1100** of FIG. **11** and may be used in the ray tracing pipeline **1000** of FIG. **10**.

[0034] In at least one embodiment, the hierarchical partitioning, such as the partitioning **120**, includes a Bounding Volume Hierarchy (BVH). However, the process **100** may be used to determine other types of partitionings of spatial elements. Examples of the spatial elements include one or more primitives, spheres, cubes, triangles, quadrilaterals, tetrahedra, points, lines, curves,

polylines, polygons, parametric shapes, polyhedrons, volumetric regions, surfaces, meshes, models, points of a point cloud, voxels, particles, Bezier patches, Non-Uniform Rational B-Spline (NURB) surfaces, implicit surfaces, fractal shapes, textured surfaces, and/or parametric surfaces. However, the partitioning **120** may include other types of data structures and/or may be used to partition other types of data. As various examples, the partitioning **120** may include one or more of an r-tree, a k-d tree, a geohash, a grid index, or a spatial hashing. Further examples of data types for the spatial elements include N-dimensional points, geographic data (e.g., geographic features, geographic coordinates or objects), sensor data (e.g., sensor readings), biological data, database data (e.g., entries, records, or elements), vector data, etc. For example, disclosed hierarchical partitionings may be used to accelerate database operations (e.g., data access) for spatially indexed data elements. While a scene is primarily described herein, the scene may also refer to a space which includes the spatial elements.

[0035] As described herein, the node manager **102** may be configured to manage nodes of the hierarchical partitioning **120** of a scene **118**, which may include modifying, creating, and/or deleting one or more nodes. The hierarchical partitioning **120** may generally include one or more nodes. For example, the hierarchical partitioning **120**A may refer to a version of the hierarchical partitioning **120** that includes a node **130**A, a node **132**A, a node **132**B, a node **134**A, a node **134**B, a node **134**C, and a node **134**D.

[0036] A node of the hierarchical partitioning **120** may include, for example, one or more values (e.g., one or more bounding shape coordinates and/or dimensions) indicating one or more bounding volumes or shapes that encloses (e.g., tightly encloses) one or more spatial elements that correspond to the node. Referring to FIG. **2** with FIG. **1**, FIG. **2** is an example of a diagram illustrating bounding shapes which may be associated with a hierarchical partitioning of spatial elements, in accordance with some embodiments of the present disclosure. FIG. **2** shows a bounding shape **230**A which may correspond to the node **130**A of the hierarchical partitioning **120**A. In the example shown, the node **130**A may correspond to a root node of the hierarchical partitioning **120**. In other examples, the node **130**A may correspond to a different type of node of the hierarchical partitioning **120**, such as an internal node or a leaf node.

[0037] In examples where the node **130**A corresponds to a root node, the node **130**A may refer to a top-level node of the hierarchical partitioning **120**, which may encompass the entire scene **118**, object, or region being partitioned. In at least one embodiment, a root node may serve as a starting point for determining the hierarchical partitioning **120** and indicate a bounding volume or shape that encloses all the spatial elements in the scene **118**. For example, the scene **118** may include any number of spatial elements, examples of which include the spatial elements **124**.

[0038] A node of the hierarchical partitioning **120** may also include, for example, one or more values (e.g., one or more pointers) indicating one or more children of the node. For example, the node **130**A may indicate the node **132**A and the node **132**B (e.g., internal nodes of the hierarchical partitioning **120**) are child nodes to the node **130**A. The children of a node may further subdivide the spatial elements **124**. For example, the child node **132**A may correspond to a bounding shape **232**A, which includes the spatial element **124**A and the spatial element **124**B. The child node **132**A may correspond to a bounding shape **232**B, which includes the spatial element **124**C and the spatial element **124**D.

[0039] Similarly, the node **132**A may indicate the node **134**A and the node **134**B (e.g., leaf nodes of the hierarchical partitioning **120**) are child nodes to the node **132**A, and the node **132**B may indicate the node **134**C and the node **134**D (e.g., leaf nodes of the hierarchical partitioning **120**) are child nodes to the node **132**B. The child node **134**A may correspond to a bounding shape **234**A, which includes the spatial element **124**A, the child node **134**B may correspond to a bounding shape **234**B, which includes the spatial element **124**B, the child node **134**C may correspond to a bounding shape **234**C, which includes the spatial element **124**C, and the child node **134**D may correspond to a bounding shape **234**D, which includes the spatial element **124**D.

[0040] A node of the hierarchical partitioning **120** may also include, for example, one or more values (e.g., one or more pointers) indicating one or more spatial elements of the node. For example, the child node **134**A indicates the spatial element **124**A, the child node **134**B may indicate the spatial element **124**B, the child node **134**C may indicate the spatial element **124**C, and the child node **134**D may indicate the spatial element **124**D. While the various nodes are described as being examples of root, internal, or leaf nodes, they may refer to other types of nodes in the process **100**. For example, the node **130**A may refer to an internal node and/or one or more of the nodes **134**A, **134**B, **134**C, or **134**D may refer to an internal node in one or more embodiments.

[0041] The nodes of the hierarchical partitioning **120** may also include addition or alternative information depending on the implementation and requirements of the application, such as, for example, bin information of multi-dimensional bins that correspond to the node (e.g., counters of spatial elements within a bin(s), aggregated bounding boxes of spatial elements within a bin(s), centroid bounds of the spatial elements within a bin(s), etc.), one or more splitting planes that correspond to the node, and/or metadata that may be used for traversal and/or modification of the hierarchical partitioning **120**.

[0042] To determine the hierarchical partitioning **120**, the process **100** may include, for example, the node manager **102** determining (e.g., generating) the node **130**A. As an example, the node **130**A may correspond to a root node of the hierarchical partitioning **120**. The bin determiner **104** may be configured to determine and/or generate multi-dimensional bins that correspond to a spatial partitioning of the scene, such as the multi-dimensional bins **122**. For example, the bin determiner **104** may be configured to determine multi-dimensional bins that correspond to a node and/or space of the scene **118** to be partitioned. In the example shown, the bin determiner **104** may be configured to determine and/or generate the multi-dimensional bins **122** corresponding to a spatial partitioning of the node **130**A and/or the bounding shape **230**A.

[0043] The multi-dimensional bins may correspond to any suitable spatial partitioning of the scene. For example, the multi-dimensional bins may include one or more grids and/or cubes enclosing regions of the scene **118**. The multi-dimensional bins **122** may, for example, correspond to a spatial partitioning of the bounding shape **230**A. The multi-dimensional bins **122** are shown as two-dimensional bins, for example, forming a grid, for simplicity. However, in at least one embodiment, the multi-dimensional bins **122** include three-dimensional bins, for example, forming a cube. Generally, a bin of the multi-dimensional bins may include at least two dimensions. The bins of the multi-dimensional bins **122** may be the same size and/or shape or one or more of the bins may vary in size along one or more dimensions. While FIG. **1** shows that the multi-dimensional bins **122** include sixteen bins, the multi-dimensional bins may include any number of bins. By way of example, for three-dimensions bins, the multi-dimensional bins may include X*Y*Z bins where X refers to the number of bins along the X-axis, Y refers to the number of bins along the Y-axis, and Z refers to the number of bins along the Z-axis. In at least one embodiment, each time the bin determiner **104** determines multi-dimensional bins for a given space to be partitioned, the multi-dimensional bins have a same quantity of bins and/or a same configuration, or different quantities or configurations may be used. For example, the multi-dimensional bins may form a multi-dimensional structure that is sized or scaled to the space to be partitioned (e.g., an 8 by 8 by 8 structure).

[0044] The bin assignor **106** may be configured to assign one or more spatial elements to one or more of the multi-dimensional bins. For example, the bin assignor **106** may be configured to assign one or more of the spatial elements **124** to one or more bins of the multi-dimensional bins **122**. In at least one embodiment, the bin assignor **106** is configured to assign a spatial element **124** to a multi-dimensional bin based at least on the bin at least partially including the spatial element. For example, the bin assignor **106** may assign each spatial element **124** to a bin based at least on determining at least one portion of the spatial element **124** falls within the bin. Referring now to FIG. **3**, FIG. **3** is an example of a diagram illustrating projection of spatial elements into multi-

dimensional bins, in accordance with some embodiments of the present disclosure.

[0045] To determine assignments between spatial elements and multi-dimensional bins, the bin assignor **106** may be configured to project spatial elements into the multi-dimensional bins along one or more dimensions. For example, one or more portions (e.g., centroids) of the spatial elements **124** are shown as being projected along an X-axis and along a Y-axis to determine which spatial elements fall within which bins along the dimensions. The bin assignor **106** may then assign the spatial elements **124** to the bins that the one or more portions fall within. For example, the spatial element **124**A may be assigned to the bin (1, 3), the spatial element **124**B may be assigned to the bin (0, 2), the spatial element **124**C may be assigned to the bin (3, 1), and the spatial element **124**D may be assigned to the bin (2, 0).

[0046] In at least one embodiment, the bin assignor **106** stores one or more values indicating the assignments between the spatial elements and the multi-dimensional bins. For example, for each bin, the bin assignor **106** may store one or more values indicating one or more spatial elements assigned to the bin. In at least one embodiment, the one or more values may indicate a count of spatial elements assigned to the one or more bins. For example, for each bin, the bin assignor **106** may store a count of spatial elements assigned to the bin.

[0047] In at least one embodiment, the one or more values may indicate an aggregation of bounding boxes of the spatial elements assigned to the one or more bins. For example, for each bin, the bin assignor **106** may store a bounding box (e.g., an X, Y, and Z coordinates) that corresponds to an aggregation of bounding boxes of spatial elements assigned to the bin.

[0048] In at least one embodiment, the one or more values may indicate the bounds of centroids or other portions of the spatial elements assigned to the one or more bins. When partitioning spatial elements using a split plane, the centroid bounds may be used to determine which side of a split plane includes the spatial element(s) from a particular bin.

[0049] Generally, the content indicated by the one or more values may vary depending upon various factors, such as the information used by the partition determiner **108** to determine partitions. In at least one embodiment, the one or more values may be stored in local memory and/or shared memory. Local memory may refer to memory that is allocated and managed at a thread or workgroup level within a PPU's execution units (e.g., streaming multiprocessors). Shared memory may refer to memory that is allocated for efficient data sharing and communication between threads within the same workgroup or thread block.

[0050] The partition determiner **108** may be configured to determine one or more partitions of the spatial elements based at least on one or more assignments between the one or more spatial elements and the one or more multi-dimensional bins determined using the bin assignor **106**. For example, the partition determiner **108** may determine, based at least on the one or more assignments and the multi-dimensional bins **122**, one or more of partitions **150**A, **150**B, **152**A, **152**B, **154**A, or **154**B.

[0051] In order for the bin assignor **106** to assign a spatial element **124** to one or more bins, location information (e.g., vertex coordinates) for the spatial element **124** may be read from memory, such as global memory (e.g., main or system memory). By assigning one or more spatial elements to one or more multi-dimensional bins, the partition determiner **108** may determine one or more partitions and one or more subpartitions thereof using a single read of spatial elements, thereby reducing the memory bandwidth needed to determine the partitions.

[0052] For example, the partition determiner **108** may determine, based at least on the one or more assignments and the multi-dimensional bins **122**, a partition, such as the partition **150**A. The partition determiner **108** may also determine, based at least on the one or more assignments and a subset of the multi-dimensional bins **122** that correspond to the partition, one or more subpartitions of the partition, such as the partition **152**A or the partition **152**B with respect to the partition **150**A. For example, the partition determiner **108** may use the one or more values indicating the assignments (e.g., counters, aggregated bounding shapes, etc.) to determine one or more partitions

and one or more sub-partitions thereof. As such, the partition determiner **108** need not determine new bins or read spatial element location information from memory to assign the spatial elements to those bins each time a partition is to be generated.

[0053] The partition determiner **108** may use various approaches to determine one or more partitions based at least on the one or more assignments and the multi-dimensional bins. In at least one embodiment, one or more partitions are determined based at least on selecting one or more split planes that define the one or more partitions. Referring now to FIG. **4**, FIG. **4** is an example of a diagram illustrating split planes which may be determined using multi-dimensional bins, in accordance with some embodiments of the present disclosure. FIG. **4** shows a split plane **410**, which the partition determiner **108** may select to define the partition **150**A and/or the partition **150**B based at least on the one or more assignments and the multi-dimensional bins **122**.

[0054] Conventionally, once a partition is determined (e.g., a split plane is selected using the bins), new bins and bin-assignments of spatial elements may be determined to determine a subpartition (e.g., a split plane is selected using the new bins). However, FIG. **4** shows a split plane **420**, which the partition determiner **108** may select to define the partition **152**A and/or the partition **152**B based at least on the one or more assignments and a set of the multi-dimensional bins **122** corresponding to the partition **150**A. FIG. **4** further shows a split plane **430**, which the partition determiner **108** may select to define the partition **154**A and/or the partition **154**B based at least on the one or more assignments and a set of the multi-dimensional bins **122** corresponding to the partition **150**B.

[0055] Various approaches may be used to select a split plane used to define a partition. In at least one embodiment, the partition determiner **108** determines one or more split plane candidates and selects one or more of the split plane candidates to define one or more partitions. For example, each set of N bins along a particular dimension may implicitly define N−1 split plane candidates where one of the N bins ends and the next bin starts. The partition determiner **108** may determine one or more such split plane candidates for one or more of the dimensions and evaluate the split plane candidates to select a split plane therefrom.

[0056] Various approaches may be used to evaluate split plane candidates. In at least one embodiment, stored bin information is used to compute one or more cost values for each split plane candidate, and the partition determiner **108** may select a split plane candidate based at least on the one or more cost values (e.g., select the lowest cost split plane candidate and/or a split plane candidate having a cost below a threshold value). For example, in at least one embodiment, the partition determiner **108** uses aggregated bin information (e.g., counts and merged bounding boxes) from any bins on each side of a candidate split plane and uses the aggregated information to compute a cost. In at least one embodiment, the cost function corresponds to a surface area of the merged bounding boxes and/or a count of the spatial elements that correspond to the merged bounding boxes. For example, the cost function may be based at least on a Surface Area Heuristic (SAH). In at least one embodiment, the cost for each split plane candidate i may be computed in accordance with Equation (1):

[00001] $\text{cost}[i] = \text{NI}[i] * \text{surfArea}(BL\_i) + \text{Nr}[i] * \text{surfArea}(Br\_i),$ (1)

where Nl[i] may refer to a sum of all bin counts for bins to the left of the split plane candidate, Nr[i] may refer to a sum of all bin counts for bins to the right of the split plane candidate, BL_i may refer to the merged bounding box of all bins to the left of the split plane candidate, BR_i may refer to the merged bounding box of all bins to the right of the split plane candidate, and surfArea may refer to a function for computing a surface area of a bounding box.

[0057] In various embodiments, any number of subpartitions (e.g., one or more) may be generated or determined from a partition using the multi-dimensional bin assignments, and any number of subpartitions may be further partitioned to generate or determine one or more additional partitions (e.g., zero or more). For example, the multi-bin assignments and a set of the multi-dimensional bins that correspond to the partition **152**B may similarly be used to determine one or more additional

split planes and partitions. However, when a partition is determined, the set of multi-dimensional bins are reduced, which may reduce the subsequent binning resolution.

[0058] In at least one embodiment, the partition determiner **108** may determine whether to partition a set of the multi-dimensional bins based at least on evaluating one or more criteria, such as one or more values indicating or representing one or more quantities of bins in the set of multi-dimensional bins along one or more dimensions (e.g., re-bin when the quantity is below a threshold for any given dimension), a quantity of candidate split planes available for partitioning, a quantity of spatial elements within the set of the multi-dimensional bins, a size of the multi-dimensional bins, a node level of the hierarchical partitioning **120**, and/or a partition or split level for the multi-dimensional bins. In at least one embodiment, based at least on evaluating the one or more criteria the partition determiner **108** may generate or determine one or more new bins for determining one or more additional partitions (e.g., generate a new set of multi-dimensional bins **122** for a region that is to be partitioned using approaches described herein).

[0059] In at least one embodiment, the projecting of the spatial elements and the evaluating of the split plane candidates may be performed warp-wide (e.g., with a thread for each bin and/or split plane) and a warp-wide reduction may be used to select a split plane from the split plane candidates. In at least one embodiment, the bin determiner **104** determines the multi-dimensional bins using a warp size. For example, where a warp includes 32 threads, the bin determiner **104** may determine a 32-bin structure that represents the multi-dimensional bins **122**. In at least one embodiment, lanes that correspond to invalid bins (e.g., indicated using shading in FIG. **4**) may be masked out (disabled or ignored) when performing projection, assignment, and split plane evaluation. While a warp is mentioned, a warp may more generally refer to a thread group, which may refer to a collection of threads that work together to solve a task and/or that execute instructions together. Examples of a thread group include a warp and a wavefront.

[0060] As described herein, the node manager **102** may be configured to assign at least one of the spatial elements to a node of the hierarchical partitioning, or otherwise update or generate the hierarchical partitioning based at least on the one or more partitions determined using the partition determiner **108**. For example, the node manager **102** may use any of the various partitions or combinations thereof to modify, create, and/or delete one or more nodes from the hierarchical partitioning **120**.

[0061] In at least one embodiment, the node manager **102** assigns one or more spatial elements that correspond to a subpartition of a partition of spatial elements to a child node to a parent node that corresponds to the partition. For example, the hierarchical partitioning **120**A may be determined using such an approach with the node **134**A corresponding to the partition **152**A and the node **134**B corresponding to the partition **152**B. The node **134**A and the node **134**B are children to the node **134**A, which corresponds to the partition **150**A. Similarly, the node **134**C corresponds to the partition **154**A and the node **134**D corresponds to the partition **154**B. The node **134**C and the node **134**D are children to the node **134**B, which corresponds to the partition **150**A.

[0062] Additionally, or in the alternative, in at least one embodiment, the node manager **102** assigns one or more spatial elements that correspond to a subpartition of a partition of spatial elements to a child node to a parent node that corresponds to the spatial elements. For example, the hierarchical partitioning **120**B may be determined using such an approach with the node **140**A corresponding to the partition **152**A, the node **140**B corresponding to the partition **152**B, the node **140**C corresponding to the partition **154**A, and the node **140**D corresponding to the partition **154**B. The node **140**A, the node **140**B, the node **140**C, and the node **140**D are children to the node **130**A. While the hierarchical partitioning **120**A and the hierarchical partitioning **120**B correspond to respectively to the approaches, hybrid approaches may be used.

[0063] After (C−1) splits have been chosen from a given set of multi-dimensional bins **122**, C child subtrees may be defined. In at least one embodiment, the generation of multi-dimensional child bins of the C children is merged into the partitioning (e.g., using merged centroid bounds, as

described herein). However, the C child bins may no longer fit into shared memory. Thus, an alternative approach may be used where during the initial construction of the multi-dimensional bins, one or more values may be stored, for each spatial element, indicating the bin coordinates for the bin assigned to the spatial element. For example, X, Y, Z coordinates may be stored where, as a non-limiting example, for 16 bins, 3*4=12 bits may be needed. The C children may then be determined (e.g., using one or more split planes), and the input spatial element IDs may be partitioned using the bin coordinates to assign each spatial element to its corresponding child partition. As such, for each spatial element, one spatial element ID and full spatial element may be read during bin construction, one bin coordinate triple may be written during construction (e.g., 12-bit), one spatial element ID and during partitioning, one triple may be read out and one spatial element ID may be written out, for a total memory footprint that may only moderately be over one full spatial element's worth of memory. The total bandwidth saved using disclosed approaches may be on the order of C/2, assuming an average of C children per multi-dimensional bin.

[0064] It may be desirable to ensure that the multi-dimensional bins to fit into shared memory. This may allow each execution unit (e.g., streaming multiprocessor) of a PPU to build its own bin over a subset of the spatial elements, and only at the end merge the bins into a global set of bins in DRAM. However, shared memory may be limited, creating a trade-off between how many bins are created (and thus, how many split planes can be selected), and how much data can be stored per bin. In at least one embodiment, to reduce the size of the data, rather than storing a full float-precision bounding box per bin, a bin may store a distance (e.g., a maximum distance) of a spatial element's bounding box coordinates from a reference position of the bin (e.g., the bin's center position). Using this approach three coordinates may be stored to indicate the bounding box rather than six. Additionally, or in the alternative, the data may be stored using lower precision, such as using 10:10:10 bit encoding, using three half-precision floats, or using three 16-bit fixed-point values relative to the spatial element bounds.

[0065] When an execution unit (e.g., streaming multiprocessor) of a PPU builds a set of bins over a set of input spatial elements, the execution unit may frequently use atomic min/max/add to update the bins. In software, the data types for representing each bin may be defied by the atomic operations. By introducing additional atomic primitives, additional data types may be available. For example, an encoded 10:10:10 three-dimensional max atomic operation may use significantly less memory per bin than an implementation using 3 floats or significantly less computationally intensive than an "emulated" 10:10:10 min atomic operation via a repeated Compare-And-Swap (CAS)-and-update atomic operation.

[0066] Now referring to FIGS. **5** and **6**, each block of method **500**, method **600**, and other methods described herein, comprises a computing process that may be performed using any combination of hardware, firmware, and/or software. For instance, various functions may be carried out by a processor executing instructions stored in memory. The methods may also be embodied as computer-usable instructions stored on computer storage media. The methods may be provided by a standalone application, a service or hosted service (standalone or in combination with another hosted service), or a plug-in to another product, to name a few. In addition, the methods are described, by way of example, with respect to the process **100** of FIG. **1**. However, these methods may additionally or alternatively be executed by any one system, or any combination of systems, including, but not limited to, those described herein.

[0067] FIG. **5** is a flow diagram showing a method **500** for projecting spatial elements into multi-dimensional bins to select split planes for partitions of spatial elements, in accordance with embodiments of the present disclosure. The method **500**, at block B**502**, includes projecting spatial elements into multi-dimensional bins to determine bin assignments (associations). For example, the bin assignor **106** may project the spatial elements **124** of the scene **118** into the multi-dimensional bins **122** that correspond to a spatial partitioning of the scene **118** to determine assignments (associations) between the spatial elements **124** and the multi-dimensional bins **122**.

[0068] At block B**504**, the method **500** includes selecting one or more first split planes using the bin assignments (associations). For example, the partition determiner **108** may select, using the assignments (associations) and the multi-dimensional bins **122**, the split plane **410** defining the partition **150**A of the spatial elements **124**.

[0069] At block B**506**, the method **500** includes selecting one or more second split planes using the bin assignments (associations). For example, the partition determiner **108** may select, using the assignments (associations) and a subset of the multi-dimensional bins **122** that is defined by the split plane **410** (e.g., the bins to the left of the split plane **410**), the split plane **420** defining the partition **152**A.

[0070] At block B**508**, the method **500** includes assigning at least one spatial element to a node. For example, the node manager **102** may assign at least one of the spatial elements **124** to the node **134**A of the hierarchical partitioning **120**A or the node **140**A of the hierarchical partitioning **120**B based at least on the partition **152**A.

[0071] FIG. **6** is a flow diagram showing a method **600** for determining values indicating assignments (associations) between spatial elements and multi-dimensional bins to select split planes for partitions of spatial elements, in accordance with embodiments of the present disclosure. The method **600**, at block B**602**, includes determining one or more values indicating assignments (associations) between spatial elements and multi-dimensional bins. For example, the bin assignor **106** may determine one or more values indicating assignments (associations) between the spatial elements **124** of the scene **118** and the multi-dimensional bins **122** that correspond to a spatial partitioning of the scene **118**.

[0072] At block B**604**, the method **600** includes selecting one or more split planes using the one or more values and the multi-dimensional bins. For example, the partition determiner **108** may select, using the one or more values and the multi-dimensional bins **122**, the split plane **410** defining the partition **150**A of the spatial elements **124**.

[0073] At block B**606**, the method **600** includes assigning at least one spatial element to a node. For example, the node manager **102** may assign at least one of the spatial elements **124** to the node **132**A of the hierarchical partitioning **120**A based at least on the partition **150**A.

EXAMPLE IMPLEMENTATIONS

[0074] A system of one or more computers can be configured to perform particular operations or actions by virtue of having software, firmware, hardware, or a combination of them installed on the system that in operation causes or cause the system to perform the actions. One or more computer programs can be configured to perform particular operations or actions by virtue of including instructions that, when executed by data processing apparatus, cause the apparatus to perform the actions. One general aspect includes a method that includes projecting spatial elements of a scene into multi-dimensional bins that correspond to a spatial partitioning of the scene to determine associations between the spatial elements and the multi-dimensional bins; selecting, using the associations and the multi-dimensional bins, one or more first split planes defining a partition of the spatial elements and a corresponding subset of the multi-dimensional bins; selecting, using the associations and the subset of the multi-dimensional bins that is defined by the one or more first split planes, one or more second split planes defining a subpartition of the partition of the spatial elements; assigning at least one spatial element of the spatial elements to a node corresponding to a hierarchical partitioning of the spatial elements based at least on the subpartition; and rendering an image of the scene by performing one or more light transport simulation techniques using the hierarchical partitioning. Other embodiments of this aspect include corresponding computer systems, apparatus, and computer programs recorded on one or more computer storage devices, each configured to perform the actions of the methods.

[0075] Implementations may include one or more of the following features. The method where the node is a child node of a parent node that corresponds to at least one of the partition or one or more of the spatial elements. The one or more first split planes are along a first dimension of the multi-

dimensional bins and the one or more second split planes are along a second dimension of the multi-dimensional bins. The selecting of the one or more first split planes is based at least on a first cost value computed for the partition using the one or more values, and the selecting of the one or more second split planes is based at least on a second cost value computed for the subpartition using the one or more values. The one or more values indicate a count of the spatial elements assigned to the one or more bins and an aggregation of bounding boxes of the spatial elements assigned to the one or more bins. A spatial element of the spatial elements includes at least one of: a primitive, a model, a mesh, a point, a voxel, a particle, a light source, an object, an emitter, or a sensor. The multi-dimensional bins include bins in at least three dimensions. The performing at least one light transport simulation technique of the one or more light transport simulation techniques may include a ray intersection testing with geometry in the scene using the hierarchical partitioning of the spatial elements as an acceleration structure. Implementations of the described techniques may include hardware, a method or process, or computer software on a computer-accessible medium.

[0076] One general aspect includes one or more processing units to perform operations including: determining one or more values indicating associations between spatial elements of a scene and multi-dimensional bins that correspond to a spatial partitioning of the scene; selecting, using the one or more values and the multi-dimensional bins, one or more split planes defining a partition of the spatial elements; assigning at least one spatial element of the spatial elements to a node corresponding to a hierarchical partitioning of the spatial elements based at least on the partition; and using the hierarchical partitioning to accelerate an application of at least one light transport simulation technique to render one or more images of the scene. Other embodiments of this aspect include corresponding computer systems, apparatus, and computer programs recorded on one or more computer storage devices, each configured to perform the actions of the methods.

[0077] Implementations may include one or more of the following features. The system where the node is a child node of a parent node that corresponds to the spatial elements. The operations further include selecting, using the associations and a subset of the multi-dimensional bins that is defined by the one or more split planes, one or more second split planes defining a subpartition of the partition of the spatial elements, where the assigning is further based at least on the subpartition. The one or more values indicate a multi-dimensional bounding box corresponding to an aggregation of bounding boxes of a plurality of the spatial elements assigned to a bin of the multi-dimensional bins. The multi-dimensional bins include at least an x-dimension, a y-dimension, and a z-dimension. The at least one light transport simulation technique may include a ray intersection testing with geometry in the scene using the hierarchical partitioning of the spatial elements as an acceleration structure. The system is may include in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing digital twin operations; a system for performing light transport simulation; a system for performing collaborative content creation for 3d assets; a system for performing deep learning operations; a system implementing one or more large language models (LLMs); a system implemented using an edge device; a system implemented using a machine; a system for performing conversational ai operations; a system for generating synthetic data; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources. Implementations of the described techniques may include hardware, a method or process, or computer software on a computer-accessible medium.

[0078] One general aspect includes at least one processor including one or more circuits to assign at least one spatial element of spatial elements of a scene to a node corresponding to a hierarchical partitioning of the spatial elements based at least on a subpartition of a partition of the spatial elements, the subpartition being determined based at least on: projecting the spatial elements into

multi-dimensional bins that correspond to a spatial partitioning of the scene to determine associations between the spatial elements and the multi-dimensional bins; determining, using the assignments and the multi-dimensional bins, the partition of the spatial elements; and determining the subpartition using the assignments and a subset of the multi-dimensional bins that corresponds to the partition. Other embodiments of this aspect include corresponding computer systems, apparatus, and computer programs recorded on one or more computer storage devices, each configured to perform the actions of the methods.

[0079] Implementations may include one or more of the following features. The at least one processor where the determining of the partition is based at least on selecting, using the assignments and the multi-dimensional bins, one or more first split planes defining the partition, and the determining of the subpartition is based at least on selecting, using the assignments and the subset of the multi-dimensional bins, one or more second split planes defining the subpartition. The node is a child node of a parent node that corresponds to at least one of the partition or one or more of the spatial elements. The one or more circuits may be to based at least on a quantity of bins in a second subset of the multi-dimensional bins that corresponds to the subpartition: project a subset of the spatial elements that correspond to subpartition into second multi-dimensional bins that correspond to a second spatial partitioning of the scene to determine second assignments between the subset of the spatial elements and the second multi-dimensional bins; determine, using the second assignments and the second multi-dimensional bins, a partition of the subset of the spatial elements; and assign one or more of the spatial elements to a second node corresponding to the hierarchical partitioning based at least on the partition of the subset of the spatial elements. The at least one processor may be comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing digital twin operations; a system for performing light transport simulation; a system for performing collaborative content creation for 3d assets; a system for performing deep learning operations; a system implementing one or more large language models (LLMs); a system implemented using an edge device; a system implemented using a machine; a system for performing conversational ai operations; a system for generating synthetic data; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources. Implementations of the described techniques may include hardware, a method or process, or computer software on a computer-accessible medium.

Example Parallel Processing Architecture

[0080] FIG. **7** illustrates an example parallel processing unit (PPU) **700** suitable for use in implementing at least some embodiments of the present disclosure. In at least one embodiment, the PPU **700** is a multi-threaded processor that is implemented on one or more integrated circuit devices. The PPU **700** may have a latency hiding architecture designed to process many threads in parallel. A thread (e.g., a thread of execution) may refer to an instantiation of a set of instructions configured to be executed by the PPU **700**. In at least one embodiment, the PPU **700** is a graphics processing unit (GPU) configured to implement a graphics rendering pipeline for processing three-dimensional (3D) graphics data in order to generate two-dimensional (2D) image data for display on a display device such as a liquid crystal display (LCD) device. In one or more embodiments, the PPU **700** may be used for performing general-purpose computations. While one parallel processor is provided herein for illustrative purposes, it should be noted that such processor is set forth for illustrative purposes only, and that any processor may be employed to supplement and/or substitute for the same.

[0081] One or more PPUs **700** may be configured to accelerate, by way of example and not limitation, thousands of High-Performance Computing (HPC), data center, and machine learning applications. The PPU **700** may be configured to accelerate numerous deep learning systems and applications including autonomous vehicle platforms, deep learning, high-accuracy speech, image,

and text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, light transport simulation, astronomy, molecular dynamics simulation, financial modeling, robotics, digital twinning, synthetic data generation, factory automation, real-time language translation, online search optimizations, personalized user recommendations, and the like.

[0082] As shown in FIG. **7**, the PPU **700** includes an Input/Output (I/O) unit **705**, a front end unit **715**, a scheduler unit **720**, a work distribution unit **725**, a hub **730**, a crossbar (Xbar) **770**, one or more general processing clusters (GPCs) **750**, and one or more partition units **780**. The PPU **700** may be connected to a host processor or other PPUs **700** via one or more high-speed NVLink **710** interconnect. The PPU **700** may be connected to a host processor or other peripheral devices via an interconnect **702**. The PPU **700** may also be connected to a local memory comprising a number of memory devices **704**. In at least one embodiment, the local memory may comprise a number of dynamic random-access memory (DRAM) devices. The DRAM devices may be configured as a high-bandwidth memory (HBM) subsystem, with multiple DRAM dies stacked within each device.

[0083] The NVLink **710** interconnect enables systems to scale and include one or more PPUs **700** combined with one or more CPUs, supports cache coherence between the PPUs **700** and CPUs, and CPU mastering. Data and/or commands may be transmitted by the NVLink **710** through the hub **730** to/from other units of the PPU **700** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown).

[0084] The I/O unit **705** may be configured to transmit and receive communications (e.g., commands, data, etc.) from a host processor (not shown) over the interconnect **702**. The I/O unit **705** may communicate with the host processor directly via the interconnect **702** or through one or more intermediate devices such as a memory bridge. In at least one embodiment, the I/O unit **705** may communicate with one or more other processors, such as one or more the PPUs **700** via the interconnect **702**. In at least one embodiment, the I/O unit **705** implements a Peripheral Component Interconnect Express (PCIe) interface for communications over a PCIe bus and the interconnect **702** is a PCIe bus. In at least one embodiment, the I/O unit **705** may implement other types of well-known interfaces for communicating with external devices.

[0085] The I/O unit **705** decodes packets received via the interconnect **702**. In at least one embodiment, the packets represent commands configured to cause the PPU **700** to perform various operations. The I/O unit **705** transmits the decoded commands to various other units of the PPU **700** as the commands may specify. For example, some commands may be transmitted to the front end unit **715**. Other commands may be transmitted to the hub **730** or other units of the PPU **700** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). In other words, the I/O unit **705** may be configured to route communications between and among the various logical units of the PPU **700**.

[0086] In at least one embodiment, a program executed by the host processor encodes a command stream in a buffer that provides workloads to the PPU **700** for processing. A workload may comprise several instructions and data to be processed by those instructions. The buffer may be a region in a memory that is accessible (e.g., read/write) by both the host processor and the PPU **700**. For example, the I/O unit **705** may be configured to access the buffer in a system memory connected to the interconnect **702** via memory requests transmitted over the interconnect **702**. In at least one embodiment, the host processor writes the command stream to the buffer and then transmits a pointer to the start of the command stream to the PPU **700**. The front end unit **715** receives pointers to one or more command streams. The front end unit **715** manages the one or more streams, reading commands from the streams and forwarding commands to the various units of the PPU **700**.

[0087] The front end unit **715** is coupled to a scheduler unit **720** that configures the various GPCs **750** to process tasks defined by the one or more streams. The scheduler unit **720** is configured to track state information related to the various tasks managed by the scheduler unit **720**. The state

may indicate which GPC **750** a task is assigned to, whether the task is active or inactive, a priority level associated with the task, and so forth. The scheduler unit **720** manages the execution of a plurality of tasks on the one or more GPCs **750**.

[0088] The scheduler unit **720** is coupled to a work distribution unit **725** that is configured to dispatch tasks for execution on the GPCs **750**. The work distribution unit **725** may track a number of scheduled tasks received from the scheduler unit **720**. In at least one embodiment, the work distribution unit **725** manages a pending task pool and an active task pool for each of the GPCs **750**. The pending task pool may comprise a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular GPC **750**. The active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by the GPCs **750**. As a GPC **750** finishes the execution of a task, that task may be evicted from the active task pool for the GPC **750** and one of the other tasks from the pending task pool is selected and scheduled for execution on the GPC **750**. If an active task has been idle on the GPC **750**, such as while waiting for a data dependency to be resolved, then the active task may be evicted from the GPC **750** and returned to the pending task pool while another task in the pending task pool is selected and scheduled for execution on the GPC **750**.

[0089] The work distribution unit **725** communicates with the one or more GPCs **750** via XBar **770**. The Xbar **770** is an interconnect network that couples many of the units of the PPU **700** to other units of the PPU **700**. For example, the Xbar **770** may be configured to couple the work distribution unit **725** to a particular GPC **750**. Although not shown explicitly, one or more other units of the PPU **700** may also be connected to the Xbar **770** via the hub **730**.

[0090] The tasks are managed by the scheduler unit **720** and dispatched to a GPC **750** by the work distribution unit **725**. The GPC **750** is configured to process the task and generate results. The results may be consumed by other tasks within the GPC **750**, routed to a different GPC **750** via the Xbar **770**, or stored in the memory **704**. The results can be written to the memory **704** via the partition units **780**, which may implement a memory interface for reading and writing data to/from the memory **704**. The results can be transmitted to another PPU **700** or CPU via the NVLink **710**. In at least one embodiment, the PPU **700** includes a number U of partition units **780** that is equal to the number of separate and distinct memory devices **704** coupled to the PPU **700**.

[0091] In at least one embodiment, a host processor executes a driver kernel that implements an application programming interface (API) that enables one or more applications executing on the host processor to schedule operations for execution on the PPU **700**. In at least one embodiment, multiple compute applications are simultaneously executed by the PPU **700** and the PPU **700** provides isolation, quality of service (QOS), and independent address spaces for the multiple compute applications. An application may generate instructions (e.g., API calls) that cause the driver kernel to generate one or more tasks for execution by the PPU **700**. The driver kernel may output tasks to one or more streams being processed by the PPU **700**. Each task may comprise one or more groups of related threads, wherein may be referred to as a warp. In at least one embodiment, a warp comprises 32 related threads that may be executed in parallel. Cooperating threads may refer to a plurality of threads including instructions to perform the task and that may exchange data through shared memory.

[0092] FIG. **8**A illustrates an example GPC **750** of the PPU **700** of FIG. **7** suitable for use in implementing at least some embodiments of the present disclosure. As shown in FIG. **8**A, each GPC **750** may include a number of hardware units for processing tasks. In at least one embodiment, each GPC **750** includes a pipeline manager **810**, a pre-raster operations unit (PROP) **815**, a raster engine **825**, a work distribution crossbar (WDX) **880**, a memory management unit (MMU) **890**, and one or more Data Processing Clusters (DPCs) **820**. It will be appreciated that the GPC **750** of FIG. **8**A may include other hardware units in lieu of or in addition to the units shown in FIG. **8**A.

[0093] In at least one embodiment, the operation of the GPC **750** is controlled by the pipeline manager **810**. The pipeline manager **810** manages the configuration of the one or more DPCs **820**

for processing tasks allocated to the GPC **750**. In at least one embodiment, the pipeline manager **810** may configure at least one of the one or more DPCs **820** to implement at least a portion of a graphics rendering pipeline. For example, a DPC **820** may be configured to execute a vertex shader program on the programmable streaming multiprocessor (SM) **840**. The pipeline manager **810** may also be configured to route packets received from the work distribution unit **725** to the appropriate logical units within the GPC **750**. For example, some packets may be routed to fixed function hardware units in the PROP **815** and/or raster engine **825** while other packets may be routed to the DPCs **820** for processing by the primitive engine **835** or the SM **840**. In at least one embodiment, the pipeline manager **810** may configure at least one of the one or more DPCs **820** to implement a neural network model and/or a computing pipeline.

[0094] The PROP unit **815** may be configured to route data generated by the raster engine **825** and the DPCs **820** to a Raster Operations (ROP) unit. The PROP unit **815** may also be configured to perform optimizations for color blending, organizing pixel data, performing address translations, and the like.

[0095] The raster engine **825** may include a number of fixed function hardware units configured to perform various raster operations. In at least one embodiment, the raster engine **825** includes a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, and a tile coalescing engine. The setup engine receives transformed vertices and generates plane equations associated with the geometric primitive defined by the vertices. The plane equations are transmitted to the coarse raster engine to generate coverage information (e.g., an x, y coverage mask for a tile) for the primitive. The output of the coarse raster engine is transmitted to the culling engine where fragments associated with the primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. Those fragments that survive clipping and culling may be passed to the fine raster engine to generate attributes for the pixel fragments based on the plane equations generated by the setup engine. The output of the raster engine **825** comprises fragments to be processed, for example, by a fragment shader implemented within a DPC **820**.

[0096] Each DPC **820** included in the GPC **750** includes an M-Pipe Controller (MPC) **830**, a primitive engine **835**, and one or more SMs **840**. The MPC **830** controls the operation of the DPC **820**, routing packets received from the pipeline manager **810** to the appropriate units in the DPC **820**. For example, packets associated with a vertex may be routed to the primitive engine **835**, which is configured to fetch vertex attributes associated with the vertex from the memory **704**. In contrast, packets associated with a shader program may be transmitted to the SM **840**.

[0097] The SM **840** comprises a programmable streaming processor that is configured to process tasks represented by a number of threads. Each SM **840** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently. In at least one embodiment, the SM **840** implements a SIMD (Single-Instruction, Multiple-Data) architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on the same set of instructions. All threads in the group of threads execute the same instructions. In at least one embodiment, the SM **840** implements a SIMT (Single-Instruction, Multiple Thread) architecture where each thread in a group of threads is configured to process a different set of data based on the same set of instructions, but where individual threads in the group of threads are allowed to diverge during execution. In at least one embodiment, a program counter, call stack, and execution state is maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within the warp diverge. In another embodiment, a program counter, call stack, and execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. When execution state is maintained for each individual thread, threads executing the same instructions may be converged and executed in parallel for maximum efficiency.

[0098] The MMU **890** may provide an interface between the GPC **750** and the partition unit **780**.

The MMU **890** may provide translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In at least one embodiment, the MMU **890** provides one or more translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in the memory **704**.

[0099] FIG. **8B** illustrates an example memory partition unit **780** of the PPU **700** of FIG. **7** suitable for use in implementing at least some embodiments of the present disclosure. As shown in FIG. **8B**, the memory partition unit **780** includes a Raster Operations (ROP) unit **850**, a level two (L2) cache **860**, and a memory interface **870**. The memory interface **870** may be coupled to the memory **704**. Memory interface **870** may implement 32, 64, 128, 1024-bit data buses, or the like, for high-speed data transfer. In at least one embodiment, the PPU **700** incorporates U memory interfaces **870**, one memory interface **870** per pair of partition units **780**, where each pair of partition units **780** is connected to a corresponding memory device **704**. For example, the PPU **700** may be connected to up to Y memory devices **704**, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random access memory, or other types of persistent storage.

[0100] In at least one embodiment, the memory interface **870** implements an HBM2 memory interface and Y equals half U. In at least one embodiment, the HBM2 memory stacks are located on the same physical package as the PPU **700**, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In at least one embodiment, each HBM2 stack includes four memory dies and Y equals 4, with HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits.

[0101] In at least one embodiment, the memory **704** supports Single-Error Correcting Double-Error Detecting (SECDED) Error Correction Code (ECC) to protect data. ECC provides high reliability for compute applications that are sensitive to data corruption. Reliability is especially important in large-scale cluster computing environments where the PPUs **700** process very large datasets and/or run applications for extended periods.

[0102] In at least one embodiment, the PPU **700** implements a multi-level memory hierarchy. In at least one embodiment, the memory partition unit **780** supports a unified memory to provide a single unified virtual address space for CPU and PPU **700** memory, enabling data sharing between virtual memory systems. In at least one embodiment the frequency of accesses by a PPU **700** to memory located on other processors is traced to ensure that memory pages are moved to the physical memory of the PPU **700** that is accessing the pages more frequently. In at least one embodiment, the NVLink **710** supports address translation services allowing the PPU **700** to directly access a CPU's page tables and providing full access to CPU memory by the PPU **700**.

[0103] In at least one embodiment, copy engines transfer data between multiple PPUs **700** or between PPUs **700** and CPUs. The copy engines can generate page faults for addresses that are not mapped into the page tables. The memory partition unit **780** can then service the page faults, mapping the addresses into the page table, after which the copy engine can perform the transfer. With hardware page faulting, addresses can be passed to the copy engines without worrying if the memory pages are resident, and the copy process is transparent.

[0104] Data from the memory **704** or other system memory may be fetched by the memory partition unit **780** and stored in the L2 cache **860**, which is located on-chip and is shared between the various GPCs **750**. As shown, each memory partition unit **780** includes a portion of the L2 cache **860** associated with a corresponding memory device **704**. Lower level caches may then be implemented in various units within the GPCs **750**. For example, each of the SMs **840** may implement a level one (L1) cache. The L1 cache is private memory that may be dedicated to a particular SM **840**. Data from the L2 cache **860** may be fetched and stored in each of the L1 caches for processing in the functional units of the SMs **840**. The L2 cache **860** is coupled to the memory interface **870** and the Xbar **770**.

[0105] The ROP unit **850** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and the like. The ROP unit **850** also implements depth testing in

conjunction with the raster engine **825**, receiving a depth for a sample location associated with a pixel fragment from the culling engine of the raster engine **825**. The depth is tested against a corresponding depth in a depth buffer for a sample location associated with the fragment. If the fragment passes the depth test for the sample location, then the ROP unit **850** updates the depth buffer and transmits a result of the depth test to the raster engine **825**. It will be appreciated that the number of partition units **780** may be different than the number of GPCs **750** and, therefore, each ROP unit **850** may be coupled to each of the GPCs **750**. The ROP unit **850** may track packets received from the different GPCs **750** and determine which GPC **750** that a result generated by the ROP unit **850** is routed to through the Xbar **770**. Although the ROP unit **850** is included within the memory partition unit **780** in FIG. **8**B, in other examples, the ROP unit **850** may be outside of the memory partition unit **780**. For example, the ROP unit **850** may reside in the GPC **750** or another unit.

[0106] FIG. **9**A illustrates an example of the streaming multi-processor **840** of FIG. **8**A suitable for use in implementing at least some embodiments of the present disclosure. As shown in FIG. **9**A, the SM **840** includes an instruction cache **905**, one or more scheduler units **912**, a register file **920**, one or more processing cores **950**, one or more special function units (SFUs) **952**, one or more load/store units (LSUs) **954**, an interconnect network **980**, and a shared memory/L1 cache **970**.

[0107] As described herein, the work distribution unit **725** dispatches tasks for execution on the GPCs **750** of the PPU **700**. The tasks may be allocated to a particular DPC **820** within a GPC **750** and, if the task is associated with a shader program, the task may be allocated to an SM **840**. The scheduler unit **912** may receive the tasks from the work distribution unit **725** and manage instruction scheduling for one or more thread blocks assigned to the SM **840**. The scheduler unit **912** may schedule thread blocks for execution as warps of parallel threads, where each thread block is allocated at least one warp. In at least one embodiment, each warp executes 32 threads. The scheduler unit **912** may manage a plurality of different thread blocks, allocating the warps to the different thread blocks and then dispatching instructions from the plurality of different cooperative groups to the various functional units (e.g., cores **950**, SFUs **952**, and LSUs **954**) during each clock cycle.

[0108] Cooperative Groups may refer to a programming model for organizing groups of communicating threads that allows developers to express the granularity at which threads are communicating, enabling the expression of richer, more efficient parallel decompositions. Cooperative launch APIs may support synchronization amongst thread blocks for the execution of parallel algorithms. Conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., the syncthreads( ) function). However, programmers would often like to define groups of threads at smaller than thread block granularities and synchronize within the defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces.

[0109] Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (e.g., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on the threads in a cooperative group. The programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

[0110] A dispatch unit **915** may be configured to transmit instructions to one or more of the functional units. In at least one embodiment, the scheduler unit **912** includes two dispatch units **915** that enable two different instructions from the same warp to be dispatched during each clock cycle. In at least embodiment, each scheduler unit **912** may include a single dispatch unit **915** or

additional dispatch units **915**.

[0111] Each SM **840** may include a register file **920** that provides a set of registers for the functional units of the SM **840**. In at least one embodiment, the register file **920** is divided between each of the functional units such that each functional unit is allocated a dedicated portion of the register file **920**. In at least one embodiment, the register file **920** is divided between the different warps being executed by the SM **840**. The register file **920** provides temporary storage for operands connected to the data paths of the functional units.

[0112] Each SM **840** may include Z processing cores **950**. In at least one embodiment, the SM **840** includes a large number (e.g., 128, etc.) of distinct processing cores **950**. Each core **950** may include a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes a floating point arithmetic logic unit and an integer arithmetic logic unit. In at least one embodiment, the floating point arithmetic logic units implement the IEEE 754-2008 standard for floating point arithmetic. In at least one embodiment, the cores **950** include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

[0113] Tensor cores configured to perform matrix operations, and, in at least one embodiment, one or more tensor cores are included in the cores **950**. In particular, the tensor cores may be configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In at least one embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation D=A×B+C, where A, B, C, and D are 4×4 matrices.

[0114] Training complex neural networks requires massive amounts of parallel computing performance, including floating-point multiplications and additions that are supported by the PPU **700**. Inferencing is less compute-intensive than training, being a latency-sensitive process where a trained neural network is applied to new inputs it has not seen before to classify images, translate speech, and infer new information.

[0115] Neural networks rely heavily on matrix math operations, and complex multi-layered networks require tremendous amounts of floating-point performance and bandwidth for both efficiency and speed. With thousands of processing cores, optimized for matrix math operations, and delivering tens to hundreds of TFLOPS of performance, the PPU **700** may form a computing platform capable of delivering performance required for deep neural network-based artificial intelligence and machine learning applications.

[0116] In at least one embodiment, the matrix multiply inputs A and B are 16-bit floating point matrices, while the accumulation matrices C and D may be 16-bit floating point or 32-bit floating point matrices. Tensor Cores operate on 16-bit floating point input data with 32-bit floating point accumulation. The 16-bit floating point multiply requires 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with the other intermediate products for a 4×4×4 matrix multiply. In practice, Tensor Cores may be used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements. An API, such as CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use Tensor Cores from a CUDA-C++ program. At the CUDA level, the warp-level interface assumes 16×16 size matrices spanning all 32 threads of the warp.

[0117] Each SM **840** may also include M SFUs **952** that perform special functions (e.g., attribute evaluation, reciprocal square root, and the like). In at least one embodiment, the SFUs **952** may include a tree traversal unit configured to traverse a hierarchical tree data structure. In at least one embodiment, the SFUs **952** may include texture unit configured to perform texture map filtering operations. In at least one embodiment, the texture units are configured to load texture maps (e.g., a 2D array of texels) from the memory **704** and sample the texture maps to produce sampled texture values for use in shader programs executed by the SM **840**. In at least one embodiment, the texture

maps are stored in the shared memory/L1 cache **870**. The texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail). In at least one embodiment, each SM **840** includes two texture units.

[0118] Each SM **840** may also include N LSUs **954** that implement load and store operations between the shared memory/L1 cache **970** and the register file **920**. Each SM **840** may include an interconnect network **980** that connects each of the functional units to the register file **920** and the LSU **954** to the register file **920**, shared memory/L1 cache **970**. In at least one embodiment, the interconnect network **980** is a crossbar that can be configured to connect any of the functional units to any of the registers in the register file **920** and connect the LSUs **954** to the register file and memory locations in shared memory/L1 cache **970**.

[0119] The shared memory/L1 cache **970** may include an array of on-chip memory that allows for data storage and communication between the SM **840** and the primitive engine **835** and between threads in the SM **840**. In at least one embodiment, the shared memory/L1 cache **970** comprises 128 KB of storage capacity and is in the path from the SM **840** to the partition unit **780**. The shared memory/L1 cache **970** can be used to cache reads and writes. One or more of the shared memory/L1 cache **970**, L2 cache **860**, and memory **704** may be backing stores.

[0120] Combining data cache and shared memory functionality into a single memory block may provide the best overall performance for both types of memory accesses. The capacity may be usable as a cache by programs that do not use shared memory. For example, if shared memory is configured to use half of the capacity, texture and load/store operations can use the remaining capacity. Integration within the shared memory/L1 cache **970** may enable the shared memory/L1 cache **970** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data.

[0121] When configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. Specifically, the fixed function graphics processing units shown in FIG. **7**, may be bypassed, creating a much simpler programming model. In the general-purpose parallel computation configuration, the work distribution unit **725** may assign and distribute blocks of threads directly to the DPCs **820**. The threads in a block may execute the same program, using a unique thread ID in the calculation to ensure each thread generates unique results, using the SM **840** to execute the program and perform calculations, shared memory/L1 cache **970** to communicate between threads, and the LSU **954** to read and write global memory through the shared memory/L1 cache **970** and the memory partition unit **780**. When configured for general purpose parallel computation, the SM **840** can also write commands that the scheduler unit **720** can use to launch new work on the DPCs **820**.

[0122] The PPU **700** may be included in a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and the like. In at least one embodiment, the PPU **700** is embodied on a single semiconductor substrate. In at least one embodiment, the PPU **700** is included in a system-on-a-chip (SoC) along with one or more other devices such as additional PPUs **700**, the memory, a reduced instruction set computer (RISC) CPU, a memory management unit (MMU), a digital-to-analog converter (DAC), and the like.

[0123] In at least one embodiment, the PPU **700** may be included on a graphics card that includes one or more memory devices **704**. The graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In at least one embodiment, the PPU **700** may be an integrated graphics processing unit (iGPU) or parallel processor included in the chipset of the motherboard.

Example of a Computing System

[0124] Systems with multiple GPUs and CPUs are used in a variety of industries as developers expose and use more parallelism in applications such as artificial intelligence computing. High-

performance GPU-accelerated systems with tens to many thousands or more of compute nodes are deployed in data centers, research facilities, and supercomputers to solve ever larger problems. As the number of processing devices within the high-performance systems increases, the communication and data transfer mechanisms need to scale to support the increased bandwidth.

[0125] FIG. **9B** is an example conceptual diagram of a processing system **900** implemented using the PPU **700** of FIG. **7** suitable for use in implementing at least some embodiments of the present disclosure. The processing system **900** includes a CPU **930**, switch **910**, and multiple PPUs **700** each and respective memories **704**. The NVLink **710** provides high-speed communication links between each of the PPUs **700**. Although a particular number of NVLink **710** and interconnect **702** connections are illustrated in FIG. **9B**, the number of connections to each PPU **700** and the CPU **930** may vary. The switch **910** interfaces between the interconnect **702** and the CPU **930**. The PPUs **700**, memories **704**, and NVLinks **710** may be situated on a single semiconductor platform to form a parallel processing system **925**. In at least one embodiment, the switch **910** supports two or more protocols to interface between various different connections and/or links.

[0126] In at least embodiment (not shown), the NVLink **710** provides one or more high-speed communication links between each of the PPUs **700** and the CPU **930** and the switch **910** interfaces between the interconnect **702** and each of the PPUs **700**. The PPUs **700**, memories **704**, and interconnect **702** may be situated on a single semiconductor platform to form a parallel processing module **925**. In at least one embodiment (not shown), the interconnect **702** provides one or more communication links between each of the PPUs **700** and the CPU **930** and the switch **910** interfaces between each of the PPUs **700** using the NVLink **710** to provide one or more high-speed communication links between the PPUs **700**. In at least one embodiment (not shown), the NVLink **710** provides one or more high-speed communication links between the PPUs **700** and the CPU **930** through the switch **910**. In yet at least one embodiment (not shown), the interconnect **702** provides one or more communication links between each of the PPUs **700** directly. One or more of the NVLink **710** high-speed communication links may be implemented as a physical NVLink interconnect or either an on-chip or on-die interconnect using the same protocol as the NVLink **710**.

[0127] In the context of the present description, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit fabricated on a die or chip. The term single semiconductor platform may also refer to multi-chip modules with increased connectivity which simulate on-chip operation and make substantial improvements over using a conventional bus implementation. Of course, the various circuits or devices may also be situated separately or in various combinations of semiconductor platforms per the desires of the user. Alternately, the parallel processing module **925** may be implemented as a circuit board substrate and each of the PPUs **700** and/or memories **704** may be packaged devices. In at least one embodiment, the CPU **930**, switch **910**, and the parallel processing module **925** are situated on a single semiconductor platform.

[0128] In at least one embodiment, the signaling rate of each NVLink **710** is 20 to 25 Gigabits/second and each PPU **700** includes six NVLink **710** interfaces (as shown in FIG. **9B**, five NVLink **710** interfaces are included for each PPU **700**). Each NVLink **710** may provide a data transfer rate of 25 Gigabytes/second in each direction, with six links providing 700 Gigabytes/second. The NVLinks **710** can be used exclusively for PPU-to-PPU communication as shown in FIG. **9B**, or some combination of PPU-to-PPU and PPU-to-CPU, when the CPU **930** also includes one or more NVLink **710** interfaces.

[0129] In at least one embodiment, the NVLink **710** allows direct load/store/atomic access from the CPU **930** to each PPU's **700** memory **704**. In at least one embodiment, the NVLink **710** supports coherency operations, allowing data read from the memories **704** to be stored in the cache hierarchy of the CPU **930**, reducing cache access latency for the CPU **930**. In at least one embodiment, the NVLink **710** includes support for Address Translation Services (ATS), allowing

the PPU **700** to directly access page tables within the CPU **930**. One or more of the NVLinks **710** may also be configured to operate in a low-power mode.

[0130] FIG. **9**C illustrates an example system **965** in which the various architecture and/or functionality of the various previous embodiments may be implemented suitable for use in implementing at least some embodiments of the present disclosure.

[0131] As shown, a system **965** is provided including at least one central processing unit **930** that is connected to a communication bus **975**. The communication bus **975** may be implemented using any suitable protocol, such as PCI (Peripheral Component Interconnect), PCI-Express, AGP (Accelerated Graphics Port), HyperTransport, or any other bus or point-to-point communication protocol(s). The system **965** also includes a main memory **940**. Control logic (software) and data are stored in the main memory **940** which may take the form of random access memory (RAM).

[0132] The system **965** also includes input devices **960**, the parallel processing system **925**, and display devices **945**, e.g. a conventional CRT (cathode ray tube), LCD (liquid crystal display), LED (light emitting diode), plasma display or the like. User input may be received from the input devices **960**, e.g., keyboard, mouse, touchpad, microphone, and the like. Each of the foregoing modules and/or devices may even be situated on a single semiconductor platform to form the system **965**. Alternately, the various modules may also be situated separately or in various combinations of semiconductor platforms per the desires of the user.

[0133] Further, the system **965** may be coupled to a network (e.g., a telecommunications network, local area network (LAN), wireless network, wide area network (WAN) such as the Internet, peer-to-peer network, cable network, or the like) through a network interface **935** for communication purposes.

[0134] The system **965** may also include a secondary storage (not shown). The secondary storage may include, for example, a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (DVD) drive, recording device, universal serial bus (USB) flash memory. The removable storage drive may read from and/or writes to a removable storage unit.

[0135] Computer programs, or computer control logic algorithms, may be stored in the main memory **940** and/or the secondary storage. Such computer programs, when executed, enable the system **965** to perform various functions. The memory **940**, the storage, and/or any other storage are possible examples of computer-readable media.

[0136] The architecture and/or functionality of the various previous figures may be implemented in the context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and/or any other desired system. For example, the system **965** may take the form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

Ray Tracing Pipeline

[0137] In at least one embodiment, the PPU **700** comprises a graphics processing unit (GPU). The PPU **700** may be configured to receive commands that specify shader programs for processing graphics data. Graphics data may be defined as a set of primitives such as points, lines, triangles, quads, triangle strips, and the like. A primitive may include data that specifies a number of vertices for the primitive (e.g., in a model-space coordinate system) as well as attributes associated with each vertex of the primitive. The PPU **700** may be configured to process the graphics primitives to generate a frame buffer (e.g., pixel data for each of the pixels of the display).

[0138] An application may write model data for a scene (e.g., a collection of vertices and attributes) to a memory such as a system memory or memory **704**. The model data may define each of the objects that may be visible on a display. The application may then make an API call to the driver

kernel that requests the model data to be rendered and displayed. The driver kernel may read the model data and write commands to the one or more streams to perform operations to process the model data. The commands may reference different shader programs to be implemented on the SMs **840** of the PPU **700**. For example, different SMs **840** may be configured to execute different shader programs.

[0139] In at least one embodiment, the model data may be processed to perform one or more ray tracing operations, such as real-time tray tracing, to render the model data to a frame buffer. The contents of the frame buffer may be transmitted to a display controller for display on a display device. Ray tracing may refer to any of a variety of techniques for modeling or simulating light transport and/or other aspects of an environment, for example, for use in generating digital images or otherwise simulating the environment. Thus, while certain embodiments may be described with respect to light transport simulation, they may be applicable to simulating, modeling, and/or measuring any of a variety of aspects of an environment. Non-limiting examples of ray tracing include ray casting, recursive ray tracing, distribution ray tracing, photon mapping, and path tracing.

[0140] Ray tracing may be used to simulate a variety of optical effects-such as shadows, reflections, refractions, scattering phenomenon, ambient occlusions, global illuminations, or dispersion phenomenon (such as chromatic aberration). Ray tracing may involve generating ray-traced samples by casting rays in a virtual environment to sample lighting and/or other environmental conditions for pixels. The ray traced samples may be combined and used to determine pixel colors for an image. In at least one embodiment, to conserve computing resources, the lighting conditions may be sparsely sampled, resulting in noisy render data. Temporal accumulation may be used to increase the effective sample count by using information from previous frames. To produce a final render that approximates a render of a fully sampled scene, one or more denoising filters may by be applied to the noisy render data to reduce noise.

[0141] Many ray tracing algorithms may cast or shoot rays from a virtual camera, or eye, through a 2D viewing plane (e.g., a pixel plane) out into a 3D scene which may include one or more light sources. Some rays may directly reach the viewing plane from a light source, some may be blocked by an object in the scene causing shadows, and some may reflect or refract off an object before reaching the viewing plane. When the rays intersect objects, the color and lighting information at the points of intersection on object surfaces may contribute to various pixel color and illumination levels of pixels of the viewing plane. Different objects may have different surface properties that can cause them to reflect, refract, or absorb light in different ways, which may be accounted for in ray tracing. Rays may reflect off objects and hit other objects, or travel through the surfaces of transparent objects before reaching a light source, and the color and lighting information from all the intersected objects may contribute to the final pixel colors.

[0142] FIG. **10** illustrates an example ray tracing pipeline **1000** suitable for use in implementing at least some embodiments of the present disclosure. By way of example, and not limitations, the ray tracing pipeline **1000** may be implemented by the PPU **700** of FIG. **7**, in accordance with at least one embodiment. The ray tracing pipeline **1000** may include processing steps implemented to generate 2D computer-generated images from 3D geometry data using one or more ray tracing techniques.

[0143] In at least one embodiment, the ray tracing pipeline **1000** may be constructed using one or more ray generation shaders **1002**, one or more any hit shaders **1004**, one or more intersection shaders **1006**, one or more miss shaders **1008**, and/or one or more closest hit shaders **1010**.

[0144] The ray tracing pipeline **1000** may be implemented via an application executed by a host processor, such as a CPU. In at least one embodiment, a device driver may implement an application programming interface (API) that defines various functions that can be used by an application in order to generate graphical data for display. The device driver may refer to a software program that includes instructions that control the operation of the PPU **700**, or other PPU used to

implement the ray tracing pipeline **1000**. The API may provide an abstraction for a programmer that lets a programmer use specialized graphics hardware, such as the PPU **700**, to generate the graphical data without requiring the programmer to use the specific instruction set for the PPU **700**. The application may include an API call that is routed to the device driver for the PPU **700**. The device driver may interpret the API call and perform various operations to respond to the API call. In at least one embodiment, the device driver performs operations by executing instructions on the CPU. In at least one embodiment, the device driver performs operations, at least in part, by launching operations on the PPU **700** using an input/output interface between the CPU and the PPU **700**. In at least one embodiment, the device driver is configured to implement the ray tracing pipeline **1000** using the hardware of the PPU **700**.

[0145] Various programs may be executed within the PPU **700** in order to implement the various stages of the ray tracing pipeline **1000**. For example, the device driver may launch a kernel on the PPU **700** to execute a stage implementing a ray generation shader **1002** on an SM **840** (or multiple SMs **840**). The device driver (or the initial kernel executed by the PPU **700**) may also launch other kernels on the PPU **700** to execute other stages of the ray tracing pipeline **1000**.

[0146] The ray generation shader **1002** may be the first shader involved in ray tracing dispatch. The ray generation shader **1002** may call a High Level Shader Language (HLSL) function called TraceRay( ). This TraceRay( ) function may cast a single ray into the scene to search for intersections, which may trigger other shaders in the process. In at least one embodiment, the ray generation shader **1002** may call TraceRay( ) any number of times.

[0147] An any hit shader **1004** and an intersection shader **1006** may be invoked whenever TraceRay( ) finds a potential intersection between the ray and the scene. The intersection shader **1006** may determine whether the ray intersects an individual geometric primitive—for example a sphere, a subdivision surface, a triangle, or other form of primitive. Once an intersection is found, the any hit shader **1004** may be used to process the intersection further or potentially discard the intersection. An any hit shader **1004** may, by way of example and not limitation, use alpha testing by performing a texture lookup and deciding based on the texel's value whether or not to discard an intersection.

[0148] Once TraceRay( ) has completed the search for ray-scene intersections, either a miss shader **1008** or a closest hit shader **1010** may be invoked, depending on the outcome of the search. The closest hit shader **1010** may perform most shading operations, such as, material evaluation, texture lookups, and so on. The miss shader **1008** may be used to implement environment lookups, for example. In at least one embodiment, one or more of the closest hit shader **1010** or the miss shader **1008** may recursively trace rays by calling TraceRay( ) themselves.

[0149] The ray tracing pipeline **1000** constructed from any of the various shaders described herein may define a single-ray programming model. In at least one embodiment, each thread of the PPU **700**, and/or other PPU used to implement the ray tracing pipeline **1000**, may handle one ray at a time. In at least one embodiment, each thread cannot communicate with other threads or see other rays currently being processed. This may simplify shader code, while allowing for vendor-specific optimizations using the API.

[0150] In at least one embodiment, different shaders and/or shader types may communicate with each other using a ray payload. A ray payload may refer to a user-defined struct that's passed as an INOUT parameter to TraceRay( ). For example, an any hit shader **1004**, a closest hit shader **1010**, and/or a miss shader **1008** may read from and/or write to the ray payload, and therefore pass back the result of their computations to the caller of TraceRay( ).

[0151] In at least one embodiment, a ray generation shader **1002** may trace primary rays, which may include rays being sent into the scene originating from a virtual camera. However, ray generation shaders **1002** are not limited to this functionality. In at least one embodiment, a ray generation shader **1002** may base ray generation on rasterized g-buffer data (e.g., to trace reflections). Using this approach, ray tracing may be used to complement rasterization, rather than

replace rasterization.

[0152] When using traditional rasterization, only the shaders required by the current object being drawn may have to be active on the PPU. This may allow rasterization pipeline objects to be relatively small, containing a single set of vertex shaders, pixel shaders, etc. In contrast, a ray tracing pipeline **1000** may be used to arbitrarily shoot rays into the scene. This may mean the rays could hit any object or many objects in the scene. Therefore, it may be the case that all shaders for all objects could potentially be hit and therefore it may be desirable for the shaders to all be resident on the PPU and ready for execution.

[0153] In at least one embodiment, a state object may be used to group shaders together for execution. At a high level, a state object of a ray tracing pipeline **1000** may be seen as a binary executable resulting from a link step across all the shaders compiled for the scene. The relationship between different shaders may be specified at state object creation. For example, triplets of intersection shaders **1006**, any hit shaders **1004**, and/or closest hit shaders **1010** may be bundled into hit groups. The application may specify the state object of the ray tracing pipeline **1000** to be executed when calling a DispatchRays( ) function on a command list. A DispathRays( ) function may invoke a ray generation shader **1002** for each pixel for an image. In at least one embodiment, an application may create any number of state objects for a ray tracing pipeline **1000** and may reuse precompiled shaders for this purpose.

[0154] Referring now to FIG. **11**, FIG. **11** illustrates an example acceleration structure **1100** suitable for use in implementing at least some embodiments of the present disclosure. The acceleration structure **1100** includes one or more top-level acceleration structures, such as a top-level acceleration structure **1102**, and one or more bottom-level acceleration structures, such as bottom-level acceleration structures **1104**A, **1104**B, and **1104**C.

[0155] The acceleration structure **1100** may comprise a spatial search data structure used in a ray tracing pipeline **1000** for acceleration structure traversal **1020** to efficiently compute intersections of rays with scene geometry. In at least one embodiment, the application may build an acceleration structure **1100** explicitly using a command list method BuildRaytracingAccelerationStructure( ). In at least one embodiment, the application may optimize an acceleration structure **1100** for different types of content, such as static versus animated content.

[0156] A top-level acceleration structure **1102** may be built from one or more references to one or more bottom-level acceleration structures **1104**A, **1104**B, and/or **1104**C. These references may be referred to as instance descriptors. Each instance descriptor may include a transformation matrix to position the instance descriptor in the scene, and an offset into a shader table **1110** (which may also be referred to as a "shader binding table") to locate material information. In at least one embodiment, a top-level acceleration structure **1102** may be used as a scene parameter provided to TraceRay( ) in a ray generation shader **1002**, and may represent an entry point of the intersection search.

[0157] A ray tracing pipeline **1000** may specify the shaders that exist in a scene and an acceleration structure **1100** may specify geometry for the scene. The shader table **1110** may refer to a data structure used to tie the geometry to the shaders. For example, the shader table **1110** may define which shader is associated with which object in the scene. In addition, the shader table **1110** may hold information about the resources accessed by each shader, such as textures, buffers, and constants.

[0158] A shader table **1110** may comprise a chunk of PPU memory, which may be managed by the application. The application may be responsible for allocating the resource, filling the shader table **1110** with valid data, transferring it to the PPU, and correctly synchronizing the shader table **1110** with ray tracing dispatches. The application may also maintain multiple shader tables **1110**, and, for example, multi-buffer them to update one copy while using another for rendering.

[0159] A shader table **1110** may comprise an array of equal-sized shader records. Each shader record may associate a shader (or a hit group) with a set of resources. In at least one embodiment,

there may exist one record per geometry object in the scene, and a shader table **1110** may include thousands of entries or more.

[0160] Referring now to FIG. **12**, FIG. **12** illustrates an example shader record **1200** suitable for use in implementing at least some embodiments of the present disclosure. The shader record **1200** is an example of a shader record that may be included in the shader table **1110** of FIG. **11**. The shader record **1200** includes a shader identifier **1202** and a root table **1204**.

[0161] In at least one embodiment, the shader identifier **1202** may be represented in a beginning portion of the shader record **1200** in memory. The shader identifier **1202** may be an opaque identifier, which the application obtains by querying for the shader identifier **1202** from a compiled shader. The root table **1204** may contain the shader's resources. The layout of the root table **1204** may be defined by the shader's local root signature. The root signature may contain any combination of constants, descriptor tables, and root descriptors. For ray tracing, the application may directly access the root table **1204** in memory (e.g., rather than using "setter" methods), which may allow for efficient updates. In at least one embodiment, a shader table **1110** may be updated from a PPU shader.

[0162] As described herein, shader table offsets may be used when building a top-level acceleration structure **1102** from instance descriptors. The system may use these offsets to locate the correct shader record **1200** whenever TraceRay( ) finds an intersection. The system may then bind the resources defined in the shader record **1200** and execute the appropriate shader for the intersected geometry.

Example Computing Device

[0163] FIG. **13** is a block diagram of an example computing device(s) **1300** suitable for use in implementing at least some embodiments of the present disclosure. Computing device **1300** may include an interconnect system **1302** that directly or indirectly couples the following devices: memory **1304**, one or more central processing units (CPUs) **1306**, one or more graphics processing units (GPUs) **1308**, a communication interface **1310**, input/output (I/O) ports **1312**, input/output components **1314**, a power supply **1316**, one or more presentation components **1318** (e.g., display(s)), and one or more logic units **1320**. In at least one embodiment, the computing device(s) **1300** may comprise one or more virtual machines (VMs), and/or any of the components thereof may comprise virtual components (e.g., virtual hardware components). For non-limiting examples, one or more of the GPUs **1308** may comprise one or more pups, one or more of the CPUs **1306** may comprise one or more vCPUs, and/or one or more of the logic units **1320** may comprise one or more virtual logic units. As such, a computing device(s) **1300** may include discrete components (e.g., a full GPU dedicated to the computing device **1300**), virtual components (e.g., a portion of a GPU dedicated to the computing device **1300**), or a combination thereof.

[0164] Although the various blocks of FIG. **13** are shown as connected via the interconnect system **1302** with lines, this is not intended to be limiting and is for clarity only. For example, in some embodiments, a presentation component **1318**, such as a display device, may be considered an I/O component **1314** (e.g., if the display is a touch screen). As another example, the CPUs **1306** and/or GPUs **1308** may include memory (e.g., the memory **1304** may be representative of a storage device in addition to the memory of the GPUs **1308**, the CPUs **1306**, and/or other components). In other words, the computing device of FIG. **13** is merely illustrative. Distinction is not made between such categories as "workstation," "server," "laptop," "desktop," "tablet," "client device," "mobile device," "hand-held device," "game console," "electronic control unit (ECU)," "virtual reality system," and/or other device or system types, as all are contemplated within the scope of the computing device of FIG. **13**.

[0165] The interconnect system **1302** may represent one or more links or busses, such as an address bus, a data bus, a control bus, or a combination thereof. The interconnect system **1302** may include one or more bus or link types, such as an industry standard architecture (ISA) bus, an extended industry standard architecture (EISA) bus, a video electronics standards association (VESA) bus, a

peripheral component interconnect (PCI) bus, a peripheral component interconnect express (PCIe) bus, and/or another type of bus or link. In some embodiments, there are direct connections between components. As an example, the CPU **1306** may be directly connected to the memory **1304**. Further, the CPU **1306** may be directly connected to the GPU **1308**. Where there is direct, or point-to-point connection between components, the interconnect system **1302** may include a PCIe link to carry out the connection. In these examples, a PCI bus need not be included in the computing device **1300**.

[0166] The memory **1304** may include any of a variety of computer-readable media. The computer-readable media may be any available media that may be accessed by the computing device **1300**. The computer-readable media may include both volatile and nonvolatile media, and removable and non-removable media. By way of example, and not limitation, the computer-readable media may comprise computer-storage media and communication media.

[0167] The computer-storage media may include both volatile and nonvolatile media and/or removable and non-removable media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, and/or other data types. For example, the memory **1304** may store computer-readable instructions (e.g., that represent a program(s) and/or a program element(s), such as an operating system. Computer-storage media may include, but is not limited to, RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical disk storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which may be used to store the desired information, and which may be accessed by computing device **1300**. As used herein, computer storage media does not comprise signals per se.

[0168] The computer storage media may embody computer-readable instructions, data structures, program modules, and/or other data types in a modulated data signal such as a carrier wave or other transport mechanism and includes any information delivery media. The term "modulated data signal" may refer to a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, the computer storage media may include wired media such as a wired network or direct-wired connection, and wireless media such as acoustic, RF, infrared and other wireless media. Combinations of any of the above should also be included within the scope of computer-readable media.

[0169] The CPU(s) **1306** may be configured to execute at least some of the computer-readable instructions to control one or more components of the computing device **1300** to perform one or more of the methods and/or processes described herein. The CPU(s) **1306** may each include one or more cores (e.g., one, two, four, eight, twenty-eight, seventy-two, etc.) that are capable of handling a multitude of software threads simultaneously. The CPU(s) **1306** may include any type of processor and may include different types of processors depending on the type of computing device **1300** implemented (e.g., processors with fewer cores for mobile devices and processors with more cores for servers). For example, depending on the type of computing device **1300**, the processor may be an Advanced RISC Machines (ARM) processor implemented using Reduced Instruction Set Computing (RISC) or an x86 processor implemented using Complex Instruction Set Computing (CISC). The computing device **1300** may include one or more CPUs **1306** in addition to one or more microprocessors or supplementary co-processors, such as math co-processors.

[0170] In addition to or alternatively from the CPU(s) **1306**, the GPU(s) **1308** may be configured to execute at least some of the computer-readable instructions to control one or more components of the computing device **1300** to perform one or more of the methods and/or processes described herein. One or more of the GPU(s) **1308** may be an integrated GPU (e.g., with one or more of the CPU(s) **1306** and/or one or more of the GPU(s) **1308** may be a discrete GPU. In embodiments, one or more of the GPU(s) **1308** may be a coprocessor of one or more of the CPU(s) **1306**. The GPU(s) **1308** may be used by the computing device **1300** to render graphics (e.g., 3D graphics) or perform

general purpose computations. For example, the GPU(s) **1308** may be used for General-Purpose computing on GPUs (GPGPU). The GPU(s) **1308** may include hundreds or thousands of cores that are capable of handling hundreds or thousands of software threads simultaneously. The GPU(s) **1308** may generate pixel data for output images in response to rendering commands (e.g., rendering commands from the CPU(s) **1306** received via a host interface). The GPU(s) **1308** may include graphics memory, such as display memory, for storing pixel data or any other suitable data, such as GPGPU data. The display memory may be included as part of the memory **1304**. The GPU(s) **1308** may include two or more GPUs operating in parallel (e.g., via a link). The link may directly connect the GPUs (e.g., using NVLINK) or may connect the GPUs through a switch (e.g., using NVSwitch). When combined together, each GPU **1308** may generate pixel data or GPGPU data for different portions of an output or for different outputs (e.g., a first GPU for a first image and a second GPU for a second image). Each GPU may include its own memory or may share memory with other GPUs.

[0171] In addition to or alternatively from the CPU(s) **1306** and/or the GPU(s) **1308**, the logic unit(s) **1320** may be configured to execute at least some of the computer-readable instructions to control one or more components of the computing device **1300** to perform one or more of the methods and/or processes described herein. In embodiments, the CPU(s) **1306**, the GPU(s) **1308**, and/or the logic unit(s) **1320** may discretely or jointly perform any combination of the methods, processes and/or portions thereof. One or more of the logic units **1320** may be part of and/or integrated in one or more of the CPU(s) **1306** and/or the GPU(s) **1308** and/or one or more of the logic units **1320** may be discrete components or otherwise external to the CPU(s) **1306** and/or the GPU(s) **1308**. In embodiments, one or more of the logic units **1320** may be a coprocessor of one or more of the CPU(s) **1306** and/or one or more of the GPU(s) **1308**.

[0172] Examples of the logic unit(s) **1320** include one or more processing cores and/or components thereof, such as Data Processing Units (DPUs), Tensor Cores (TCs), Tensor Processing Units (TPUs), Pixel Visual Cores (PVCs), Vision Processing Units (VPUs), Graphics Processing Clusters (GPCs), Texture Processing Clusters (TPCs), Streaming Multiprocessors (SMs), Tree Traversal Units (TTUs), Artificial Intelligence Accelerators (AIAs), Deep Learning Accelerators (DLAs), Arithmetic-Logic Units (ALUs), Application-Specific Integrated Circuits (ASICs), Floating Point Units (FPUs), input/output (I/O) elements, peripheral component interconnect (PCI) or peripheral component interconnect express (PCIe) elements, and/or the like.

[0173] The communication interface **1310** may include one or more receivers, transmitters, and/or transceivers that enable the computing device **1300** to communicate with other computing devices via an electronic communication network, included wired and/or wireless communications. The communication interface **1310** may include components and functionality to enable communication over any of a number of different networks, such as wireless networks (e.g., Wi-Fi, Z-Wave, Bluetooth, Bluetooth LE, ZigBee, etc.), wired networks (e.g., communicating over Ethernet or InfiniBand), low-power wide-area networks (e.g., LoRaWAN, SigFox, etc.), and/or the Internet. In one or more embodiments, logic unit(s) **1320** and/or communication interface **1310** may include one or more data processing units (DPUs) to transmit data received over a network and/or through interconnect system **1302** directly to (e.g., a memory of) one or more GPU(s) **1308**.

[0174] The I/O ports **1312** may enable the computing device **1300** to be logically coupled to other devices including the I/O components **1314**, the presentation component(s) **1318**, and/or other components, some of which may be built in to (e.g., integrated in) the computing device **1300**. Illustrative I/O components **1314** include a microphone, mouse, keyboard, joystick, game pad, game controller, satellite dish, scanner, printer, wireless device, etc. The I/O components **1314** may provide a natural user interface (NUI) that processes air gestures, voice, or other physiological inputs generated by a user. In some instances, inputs may be transmitted to an appropriate network element for further processing. An NUI may implement any combination of speech recognition, stylus recognition, facial recognition, biometric recognition, gesture recognition both on screen and

adjacent to the screen, air gestures, head and eye tracking, and touch recognition (as described in more detail below) associated with a display of the computing device **1300**. The computing device **1300** may be include depth cameras, such as stereoscopic camera systems, infrared camera systems, RGB camera systems, touchscreen technology, and combinations of these, for gesture detection and recognition. Additionally, the computing device **1300** may include accelerometers or gyroscopes (e.g., as part of an inertia measurement unit (IMU)) that enable detection of motion. In some examples, the output of the accelerometers or gyroscopes may be used by the computing device **1300** to render immersive augmented reality or virtual reality.

[0175] The power supply **1316** may include a hard-wired power supply, a battery power supply, or a combination thereof. The power supply **1316** may provide power to the computing device **1300** to enable the components of the computing device **1300** to operate.

[0176] The presentation component(s) **1318** may include a display (e.g., a monitor, a touch screen, a television screen, a heads-up-display (HUD), other display types, or a combination thereof), speakers, and/or other presentation components. The presentation component(s) **1318** may receive data from other components (e.g., the GPU(s) **1308**, the CPU(s) **1306**, DPUs, etc.), and output the data (e.g., as an image, video, sound, etc.).

Example Data Center

[0177] FIG. **14** illustrates an example data center **1400** that may be used in at least one embodiments of the present disclosure. The data center **1400** may include a data center infrastructure layer **1410**, a framework layer **1420**, a software layer **1430**, and/or an application layer **1440**.

[0178] As shown in FIG. **14**, the data center infrastructure layer **1410** may include a resource orchestrator **1412**, grouped computing resources **1414**, and node computing resources ("node C.R.s") **1416(1)**-**1416**(N), where "N" represents any whole, positive integer. In at least one embodiment, node C.R.s **1416(1)**-**1416**(N) may include, but are not limited to, any number of central processing units (CPUs) or other processors (including DPUs, accelerators, field programmable gate arrays (FPGAs), graphics processors or graphics processing units (GPUs), etc.), memory devices (e.g., dynamic read-only memory), storage devices (e.g., solid state or disk drives), network input/output (NW I/O) devices, network switches, virtual machines (VMs), power modules, and/or cooling modules, etc. In some embodiments, one or more node C.R.s from among node C.R.s **1416(1)**-**1416**(N) may correspond to a server having one or more of the above-mentioned computing resources. In addition, in some embodiments, the node C.R.s **1416(1)**-**14161** (N) may include one or more virtual components, such as vGPUs, vCPUs, and/or the like, and/or one or more of the node C.R.s **1416(1)**-**1416**(N) may correspond to a virtual machine (VM).

[0179] In at least one embodiment, grouped computing resources **1414** may include separate groupings of node C.R.s **1416** housed within one or more racks (not shown), or many racks housed in data centers at various geographical locations (also not shown). Separate groupings of node C.R.s **1416** within grouped computing resources **1414** may include grouped compute, network, memory or storage resources that may be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s **1416** including CPUs, GPUs, DPUs, and/or other processors may be grouped within one or more racks to provide compute resources to support one or more workloads. The one or more racks may also include any number of power modules, cooling modules, and/or network switches, in any combination.

[0180] The resource orchestrator **1412** may configure or otherwise control one or more node C.R.s **1416(1)**-**1416**(N) and/or grouped computing resources **1414**. In at least one embodiment, resource orchestrator **1412** may include a software design infrastructure (SDI) management entity for the data center **1400**. The resource orchestrator **1412** may include hardware, software, or some combination thereof.

[0181] In at least one embodiment, as shown in FIG. **14**, framework layer **1420** may include a job scheduler **1428**, a configuration manager **1434**, a resource manager **1436**, and/or a distributed file

system **1438**. The framework layer **1420** may include a framework to support software **1432** of software layer **1430** and/or one or more application(s) **1442** of application layer **1440**. The software **1432** or application(s) **1442** may respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud and Microsoft Azure. The framework layer **1420** may be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter "Spark") that may utilize distributed file system **1438** for large-scale data processing (e.g., "big data"). In at least one embodiment, job scheduler **1428** may include a Spark driver to facilitate scheduling of workloads supported by various layers of data center **1400**. The configuration manager **1434** may be capable of configuring different layers such as software layer **1430** and framework layer **1420** including Spark and distributed file system **1438** for supporting large-scale data processing. The resource manager **1436** may be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system **1438** and job scheduler **1428**. In at least one embodiment, clustered or grouped computing resources may include grouped computing resource **1414** at data center infrastructure layer **1410**. The resource manager **1436** may coordinate with resource orchestrator **1412** to manage these mapped or allocated computing resources.

[0182] In at least one embodiment, software **1432** included in software layer **1430** may include software used by at least portions of node C.R.s **1416(1)**-**1416**(N), grouped computing resources **1414**, and/or distributed file system **1438** of framework layer **1420**. One or more types of software may include, but are not limited to, Internet web page search software, e-mail virus scan software, database software, and streaming video content software.

[0183] In at least one embodiment, application(s) **1442** included in application layer **1440** may include one or more types of applications used by at least portions of node C.R.s **1416(1)**-**1416**(N), grouped computing resources **1414**, and/or distributed file system **1438** of framework layer **1420**. One or more types of applications may include, but are not limited to, any number of a genomics application, a cognitive compute, and a machine learning application, including training or inferencing software, machine learning framework software (e.g., PyTorch, TensorFlow, Caffe, etc.), and/or other machine learning applications used in conjunction with one or more embodiments.

[0184] In at least one embodiment, any of configuration manager **1434**, resource manager **1436**, and resource orchestrator **1412** may implement any number and type of self-modifying actions based on any amount and type of data acquired in any technically feasible fashion. Self-modifying actions may relieve a data center operator of data center **1400** from making possibly bad configuration decisions and possibly avoiding underused and/or poor performing portions of a data center.

[0185] The data center **1400** may include tools, services, software or other resources to train one or more machine learning models or predict or infer information using one or more machine learning models according to one or more embodiments described herein. For example, a machine learning model(s) may be trained by calculating weight parameters according to a neural network architecture using software and/or computing resources described above with respect to the data center **1400**. In at least one embodiment, trained or deployed machine learning models corresponding to one or more neural networks may be used to infer or predict information using resources described above with respect to the data center **1400** by using weight parameters calculated through one or more training techniques, such as but not limited to those described herein.

[0186] In at least one embodiment, the data center **1400** may use CPUs, application-specific integrated circuits (ASICs), GPUs, FPGAs, and/or other hardware (or virtual compute resources corresponding thereto) to perform training and/or inferencing using above-described resources. Moreover, one or more software and/or hardware resources described above may be configured as a service to allow users to train or performing inferencing of information, such as image

recognition, speech recognition, or other artificial intelligence services.

Example Network Environments

[0187] Network environments suitable for use in implementing embodiments of the disclosure may include one or more client devices, servers, network attached storage (NAS), other backend devices, and/or other device types. The client devices, servers, and/or other device types (e.g., each device) may be implemented on one or more instances of the computing device(s) **1300** of FIG. **13** —e.g., each device may include similar components, features, and/or functionality of the computing device(s) **1300**. In addition, where backend devices (e.g., servers, NAS, etc.) are implemented, the backend devices may be included as part of a data center **1400**, an example of which is described in more detail herein with respect to FIG. **14**.

[0188] Components of a network environment may communicate with each other via a network(s), which may be wired, wireless, or both. The network may include multiple networks, or a network of networks. By way of example, the network may include one or more Wide Area Networks (WANs), one or more Local Area Networks (LANs), one or more public networks such as the Internet and/or a public switched telephone network (PSTN), and/or one or more private networks. Where the network includes a wireless telecommunications network, components such as a base station, a communications tower, or even access points (as well as other components) may provide wireless connectivity.

[0189] Compatible network environments may include one or more peer-to-peer network environments—in which case a server may not be included in a network environment—and one or more client-server network environments—in which case one or more servers may be included in a network environment. In peer-to-peer network environments, functionality described herein with respect to a server(s) may be implemented on any number of client devices.

[0190] In at least one embodiment, a network environment may include one or more cloud-based network environments, a distributed computing environment, a combination thereof, etc. A cloud-based network environment may include a framework layer, a job scheduler, a resource manager, and a distributed file system implemented on one or more of servers, which may include one or more core network servers and/or edge servers. A framework layer may include a framework to support software of a software layer and/or one or more application(s) of an application layer. The software or application(s) may respectively include web-based service software or applications. In embodiments, one or more of the client devices may use the web-based service software or applications (e.g., by accessing the service software and/or applications via one or more application programming interfaces (APIs)). The framework layer may be, but is not limited to, a type of free and open-source software web application framework such as that may use a distributed file system for large-scale data processing (e.g., "big data").

[0191] A cloud-based network environment may provide cloud computing and/or cloud storage that carries out any combination of computing and/or data storage functions described herein (or one or more portions thereof). Any of these various functions may be distributed over multiple locations from central or core servers (e.g., of one or more data centers that may be distributed across a state, a region, a country, the globe, etc.). If a connection to a user (e.g., a client device) is relatively close to an edge server(s), a core server(s) may designate at least a portion of the functionality to the edge server(s). A cloud-based network environment may be private (e.g., limited to a single organization), may be public (e.g., available to many organizations), and/or a combination thereof (e.g., a hybrid cloud environment).

[0192] The client device(s) may include at least some of the components, features, and functionality of the example computing device(s) **1300** described herein with respect to FIG. **13**. By way of example and not limitation, a client device may be embodied as a Personal Computer (PC), a laptop computer, a mobile device, a smartphone, a tablet computer, a smart watch, a wearable computer, a Personal Digital Assistant (PDA), an MP3 player, a virtual reality headset, a Global Positioning System (GPS) or device, a video player, a video camera, a surveillance device

or system, a vehicle, a boat, a flying vessel, a virtual machine, a drone, a robot, a handheld communications device, a hospital device, a gaming device or system, an entertainment system, a vehicle computer system, an embedded system controller, a remote control, an appliance, a consumer electronic device, a workstation, an edge device, any combination of these delineated devices, or any other suitable device.

[0193] The disclosure may be described in the general context of computer code or machine-useable instructions, including computer-executable instructions such as program modules, being executed by a computer or other machine, such as a personal data assistant or other handheld device. Generally, program modules including routines, programs, objects, components, data structures, etc., refer to code that perform particular tasks or implement particular abstract data types. The disclosure may be practiced in a variety of system configurations, including hand-held devices, consumer electronics, general-purpose computers, more specialty computing devices, etc. The disclosure may also be practiced in distributed computing environments where tasks are performed by remote-processing devices that are linked through a communications network.

[0194] As used herein, a recitation of "and/or" with respect to two or more elements should be interpreted to mean only one element, or a combination of elements. For example, "element A, element B, and/or element C" may include only element A, only element B, only element C, element A and element B, element A and element C, element B and element C, or elements A, B, and C. In addition, "at least one of element A or element B" may include at least one of element A, at least one of element B, or at least one of element A and at least one of element B. Further, "at least one of element A and element B" may include at least one of element A, at least one of element B, or at least one of element A and at least one of element B.

[0195] The subject matter of the present disclosure is described with specificity herein to meet statutory requirements. However, the description itself is not intended to limit the scope of this disclosure. Rather, the inventors have contemplated that the claimed subject matter might also be embodied in other ways, to include different steps or combinations of steps similar to the ones described in this document, in conjunction with other present or future technologies. Moreover, although the terms "step" and/or "block" may be used herein to connote different elements of methods employed, the terms should not be interpreted as implying any particular order among or between various steps herein disclosed unless and except when the order of individual steps is explicitly described.

Claims

1. A method comprising: projecting spatial elements of a scene into multi-dimensional bins that correspond to a spatial partitioning of the scene to determine associations between the spatial elements and the multi-dimensional bins; selecting, using the associations and the multi-dimensional bins, one or more first split planes defining a partition of the spatial elements and a corresponding subset of the multi-dimensional bins; selecting, using the associations and the subset of the multi-dimensional bins that is defined by the one or more first split planes, one or more second split planes defining a subpartition of the partition of the spatial elements; assigning at least one spatial element of the spatial elements to a node corresponding to a hierarchical partitioning of the spatial elements based at least on the subpartition; and rendering an image of the scene by performing one or more light transport simulation techniques using the hierarchical partitioning.

2. The method of claim 1, wherein the node is a child node of a parent node that corresponds to at least one of the partition or one or more of the spatial elements.

3. The method of claim 1, wherein the one or more first split planes are along a first dimension of the multi-dimensional bins and the one or more second split planes are along a second dimension of the multi-dimensional bins.

4. The method of claim 1, further comprising computing one or more values indicating the

associations with respect to one or more bins of the multi-dimensional bins, wherein the selecting of the one or more first split planes is based at least on a first cost value computed for the partition using the one or more values, and the selecting of the one or more second split planes is based at least on a second cost value computed for the subpartition using the one or more values.

5. The method of claim 4, wherein the one or more values indicate a count of the spatial elements assigned to the one or more bins and an aggregation of bounding boxes of the spatial elements assigned to the one or more bins.

6. The method of claim 1, wherein a spatial element of the spatial elements includes at least one of: a primitive, a model, a mesh, a point, a voxel, a particle, a light source, an object, an emitter, or a sensor.

7. The method of claim 1, wherein the multi-dimensional bins include bins in at least three dimensions.

8. The method of claim 1, wherein the performing at least one light transport simulation technique of the one or more light transport simulation techniques comprises a ray intersection testing with geometry in the scene using the hierarchical partitioning of the spatial elements as an acceleration structure.

9. A system comprising: one or more processing units to perform operations including: determining one or more values indicating associations between spatial elements of a scene and multi-dimensional bins that correspond to a spatial partitioning of the scene; selecting, using the one or more values and the multi-dimensional bins, one or more split planes defining a partition of the spatial elements; assigning at least one spatial element of the spatial elements to a node corresponding to a hierarchical partitioning of the spatial elements based at least on the partition; and using the hierarchical partitioning to accelerate an application of at least one light transport simulation technique to render one or more images of the scene.

10. The system of claim 9, wherein the node is a child node of a parent node that corresponds to the spatial elements.

11. The system of claim 9, wherein the operations further include selecting, using the associations and a subset of the multi-dimensional bins that is defined by the one or more split planes, one or more second split planes defining a subpartition of the partition of the spatial elements, wherein the assigning is further based at least on the subpartition.

12. The system of claim 9, wherein the one or more values indicate a multi-dimensional bounding box corresponding to an aggregation of bounding boxes of a plurality of the spatial elements assigned to a bin of the multi-dimensional bins.

13. The system of claim 9, wherein the multi-dimensional bins include at least an x-dimension, a y-dimension, and a z-dimension.

14. The system of claim 9, wherein the at least one light transport simulation technique comprises a ray intersection testing with geometry in the scene using the hierarchical partitioning of the spatial elements as an acceleration structure.

15. The system of claim 9, wherein the system is comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing digital twin operations; a system for performing light transport simulation; a system for performing collaborative content creation for 3D assets; a system for performing deep learning operations; a system implementing one or more large language models (LLMs); a system implemented using an edge device; a system implemented using a machine; a system for performing conversational AI operations; a system for generating synthetic data; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources.

16. A processor comprising: one or more circuits to assign at least one spatial element of spatial elements of a scene to a node corresponding to a hierarchical partitioning of the spatial elements

based at least on a subpartition of a partition of the spatial elements, the subpartition being determined based at least on: projecting the spatial elements into multi-dimensional bins that correspond to a spatial partitioning of the scene to determine associations between the spatial elements and the multi-dimensional bins; determining, using the assignments and the multi-dimensional bins, the partition of the spatial elements; and determining the subpartition using the assignments and a subset of the multi-dimensional bins that corresponds to the partition.

17. The processor of claim 16, wherein the determining of the partition is based at least on selecting, using the assignments and the multi-dimensional bins, one or more first split planes defining the partition, and the determining of the subpartition is based at least on selecting, using the assignments and the subset of the multi-dimensional bins, one or more second split planes defining the subpartition.

18. The processor of claim 16, wherein the node is a child node of a parent node that corresponds to at least one of the partition or one or more of the spatial elements.

19. The processor of claim 16, wherein the one or more circuits are to based at least on a quantity of bins in a second subset of the multi-dimensional bins that corresponds to the subpartition: project a subset of the spatial elements that correspond to subpartition into second multi-dimensional bins that correspond to a second spatial partitioning of the scene to determine second assignments between the subset of the spatial elements and the second multi-dimensional bins; determine, using the second assignments and the second multi-dimensional bins, a partition of the subset of the spatial elements; and assign one or more of the spatial elements to a second node corresponding to the hierarchical partitioning based at least on the partition of the subset of the spatial elements.

20. The processor of claim 16, wherein the processor is comprised in at least one of: a control system for an autonomous or semi-autonomous machine; a perception system for an autonomous or semi-autonomous machine; a system for performing simulation operations; a system for performing digital twin operations; a system for performing light transport simulation; a system for performing collaborative content creation for 3D assets; a system for performing deep learning operations; a system implementing one or more large language models (LLMs); a system implemented using an edge device; a system implemented using a machine; a system for performing conversational AI operations; a system for generating synthetic data; a system incorporating one or more virtual machines (VMs); a system implemented at least partially in a data center; or a system implemented at least partially using cloud computing resources.