



US 20250259403A1

(19) **United States**

(12) **Patent Application Publication**

Velasquez et al.

(10) **Pub. No.: US 2025/0259403 A1**

(43) **Pub. Date:** Aug. 14, 2025

(54) **CROSS REALITY SYSTEM WITH ACCURATE SHARED MAPS**

(71) Applicant: **Magic Leap, Inc.**, Plantation, FL (US)

(72) Inventors: **Miguel Andres Granados Velasquez**, Thalwil (CH); **Javier Victorio Gomez Gonzalez**, Zurich (CH); **Mukta Prasad**, Uitikon Waldegg (CH); **Eran Guendelman**, Tel Aviv (IL); **Ali Shahroknii**, San Jose, CA (US); **Ashwin Swaminathan**, Dublin, CA (US)

(73) Assignee: **Magic Leap, Inc.**, Plantation, FL (US)

(21) Appl. No.: **19/169,840**

(22) Filed: **Apr. 3, 2025**

Related U.S. Application Data

(63) Continuation of application No. 18/457,314, filed on Aug. 28, 2023, now Pat. No. 12,283,011, which is a continuation of application No. 17/856,903, filed on Jul. 1, 2022, now Pat. No. 11,790,619, which is a continuation of application No. 17/174,122, filed on Feb. 11, 2021, now Pat. No. 11,410,395.

(60) Provisional application No. 62/975,983, filed on Feb. 13, 2020.

Publication Classification

(51) **Int. Cl.**

G06T 19/00 (2011.01)

G01C 21/00 (2006.01)


(52) **U.S. Cl.**

CPC **G06T 19/006** (2013.01); **G01C 21/387** (2020.08)

(57)

ABSTRACT

A cross reality system enables any of multiple devices to efficiently and accurately access previously persisted maps of very large scale environments and render virtual content specified in relation to those maps. The cross reality system may build a persisted map, which may be in canonical form, by merging tracking maps from the multiple devices. A map merge process determines mergibility of a tracking map with a canonical map and merges a tracking map with a canonical map in accordance with mergibility criteria, such as, when a gravity direction of the tracking map aligns with a gravity direction of the canonical map. Refraining from merging maps if the orientation of the tracking map with respect to gravity is not preserved avoids distortions in persisted maps and results in multiple devices, which may use the maps to determine their locations, to present more realistic and immersive experiences for their users.






FIG. 1




FIG. 2




FIG. 3




FIG. 4




FIG. 5A




FIG. 5B





FIG. 6A







FIG. 7






FIG. 9




FIG. 10




FIG. 11




FIG. 12




FIG. 13




FIG. 14





FIG. 15






FIG. 17




FIG. 18

FIG. 19




FIG. 20A




FIG. 20B




FIG. 20C




FIG. 21




FIG. 22




FIG. 23




FIG. 24




FIG. 25




FIG. 26




FIG. 27




FIG. 28




FIG. 29




FIG. 30

FIG. 31B





FIG. 31A






FIG. 32




FIG. 33





FIG. 34




FIG. 35




FIG. 36




FIG. 37




FIG. 38A







FIG. 38B






FIG. 39D




FIG. 39E




FIG. 39F




FIG. 40




FIG. 41




FIG. 42




FIG. 43A




FIG. 43B




FIG. 43C

Second XR Device 12.2




FIG. 44





FIG. 45



Second XR Device 12.2





FIG. 46B




FIG. 47




FIG. 48




FIG. 49




FIG. 50




FIG. 51

Second XR Device 12.2

Map 2 Localized
To New Canonical
Map




FIG. 52

Second XR Device 12.2




FIG. 53




FIG. 54




FIG. 55




FIG. 56




FIG. 57




FIG. 58




FIG. 59




FIG. 60




FIG. 61




FIG. 62




FIG. 63A




FIG. 63B




FIG. 63C




FIG. 64




FIG. 65




FIG. 66




FIG. 67




FIG. 68




FIG. 69

CROSS REALITY SYSTEM WITH ACCURATE SHARED MAPS

CROSS-REFERENCE TO RELATED APPLICATIONS

[0001] This application is a continuation of U.S. patent application Ser. No. 18/457,314, filed on Aug. 28, 2023, entitled “CROSS REALITY SYSTEM WITH ACCURATE SHARED MAPS,” which is a continuation of U.S. patent application Ser. No. 17/856,903, now U.S. Pat. No. 11,790,619, filed on Jul. 1, 2022, entitled “CROSS REALITY SYSTEM WITH ACCURATE SHARED MAPS,” which is a continuation of U.S. patent application Ser. No. 17/174,122, now U.S. Pat. No. 11,410,395, filed on Feb. 11, 2021, entitled “CROSS REALITY SYSTEM WITH ACCURATE SHARED MAPS,” which claims priority to and the benefit under 35 U.S.C. § 119(e) of U.S. Provisional Patent Application Ser. No. 62/975,983, filed on Feb. 13, 2020, entitled “CROSS REALITY SYSTEM WITH ACCURATE SHARED MAPS.” The contents of these applications are hereby incorporated herein by reference in their entirety.

TECHNICAL FIELD

[0002] This application relates generally to a cross reality system.

BACKGROUND

[0003] Computers may control human user interfaces to create a cross reality (XR) environment in which some or all of the XR environment, as perceived by the user, is generated by the computer. These XR environments may be virtual reality (VR), augmented reality (AR), and mixed reality (MR) environments, in which some or all of an XR environment may be generated by computers using, in part, data that describes the environment. This data may describe, for example, virtual objects that may be rendered in a way that users’ sense or perceive as a part of a physical world and can interact with the virtual objects. The user may experience these virtual objects as a result of the data being rendered and presented through a user interface device, such as, for example, a head-mounted display device. The data may be displayed to the user to see, or may control audio that is played for the user to hear, or may control a tactile (or haptic) interface, enabling the user to experience touch sensations that the user senses or perceives as feeling the virtual object.

[0004] XR systems may be useful for many applications, spanning the fields of scientific visualization, medical training, engineering design and prototyping, tele-manipulation and tele-presence, and personal entertainment. AR and MR, in contrast to VR, include one or more virtual objects in relation to real objects of the physical world. The experience of virtual objects interacting with real objects greatly enhances the user’s enjoyment in using the XR system, and also opens the door for a variety of applications that present realistic and readily understandable information about how the physical world might be altered.

[0005] To realistically render virtual content, an XR system may build a representation of the physical world around a user of the system. This representation, for example, may be constructed by processing images acquired with sensors on a wearable device that forms a part of the XR system. In such a system, a user might perform an initialization routine

by looking around a room or other physical environment in which the user intends to use the XR system until the system acquires sufficient information to construct a representation of that environment. As the system operates and the user moves around the environment or to other environments, the sensors on the wearable devices might acquire additional information to expand or update the representation of the physical world.

BRIEF SUMMARY

[0006] Aspects of the present application relate to methods and apparatus for providing cross reality (XR) scenes. Techniques as described herein may be used together, separately, or in any suitable combination.

[0007] According to one embodiment, a method of merging one or more environment maps stored in a database with a tracking map computed based on sensor data collected by a device worn by a user is provided. The method may comprise receiving the tracking map from the device, wherein the tracking map is aligned with respect to a gravity direction; determining a transformation between the tracking map and an environment map; determining whether to merge the environment map with the tracking map, wherein determining whether to merge comprises determining whether applying the transformation to the tracking map produces a transformed tracking map that aligns with respect to the gravity direction; and merging the environment map with the tracking map based on determining that the transformed tracking map aligns with the gravity direction.

[0008] According to one embodiment, determining the transformation may comprise, for corresponding features in the tracking map and the environment map, selecting as the determined transformation, the transformation that aligns with the corresponding features with a metric of error below a threshold.

[0009] According to one embodiment, the method may further comprising determining the corresponding features based on similarity of identifiers assigned to the features.

[0010] According to one embodiment, selecting as the determined transformation may further comprise selecting the transformation that when applied produces the transformed tracking map that aligns with respect to the gravity direction.

[0011] According to one embodiment, determining the transformation may comprise applying a plurality of candidate transformations to the tracking map and selecting as the determined transformation a candidate transformation of the plurality of candidate transformations.

[0012] According to one embodiment, selecting the determined transformation may further comprise selecting as the determined transformation the candidate transformation that when applied produces the transformed tracking map that aligns with respect to the gravity direction.

[0013] According to one embodiment, determining whether applying the transformation to the tracking map produces the transformed tracking map that aligns with respect to the gravity direction comprises determining whether applying the transformation to the tracking map produces the transformed tracking map that is rotated with respect to the gravity direction by more than a threshold amount; in response to determining that applying the transformation to the tracking map does not produce the transformed tracking map that is rotated with respect to the gravity direction by more than the threshold amount, select-

ing the transformation as the determined transformation; and in response to determining that applying the transformation to the tracking map produces the transformed tracking map that is rotated with respect to the gravity direction by more than the threshold amount, discarding the transformation.

[0014] According to one embodiment, the method may further comprise identifying a set of environment maps from the database to be merged with the tracking map; and for each environment map in the set of environment maps: determining the transformation between the tracking map and the environment map; determining whether to merge the environment map with the tracking map; and merging the environment map with the tracking map based on determining that the transformed tracking map aligns with the gravity direction.

[0015] According to one embodiment, the method may further comprise for each environment map in the set of environment maps: refraining from merging the environment map with the tracking map based on determining that a gravity direction of the environment map does not align with the gravity direction of the transformed tracking map.

[0016] According to one embodiment, identifying the set of environment maps may comprise determining area identifiers associated with the tracking map; and identifying the set of environment maps from the database based, at least in part on, on the area identifiers associated with the tracking map.

[0017] According to one embodiment, identifying the set of environment maps from the database may further comprise filtering the set of environment maps based on similarity of one or more metrics associated with the tracking map and the environment maps in the set of environment maps.

[0018] According to one embodiment, a computing device is configured for use in a cross reality system in which a portable device operating in a three-dimensional (3D) environment renders virtual content. The computing device may comprise at least one processor; a computer-readable medium connected to the processor; a plurality of environment maps stored in the computer-readable medium; and computer-executable instructions configured to, when executed by the at least one processor, perform a method. The method performed by the computer-executable instructions may comprise receiving a tracking map from the portable device, wherein the tracking map is aligned with respect to a gravity direction; determining whether to merge an environment map with the tracking map, wherein determining whether to merge comprises searching for a transformation of the tracking map that aligns a transformed tracking map and the environment map in a manner that preserves alignment of the transformed tracking map with respect to the gravity direction; and merging the environment map with the transformed tracking map based on determining that a gravity direction of the environment map aligns with the gravity direction of the transformed tracking map.

[0019] According to one embodiment, searching for a transformation may comprise searching for the transformation that aligns a first set of features associated with the tracking map with a second set of features associated with the environment map with a metric of error below a threshold.

[0020] According to one embodiment, searching for a transformation may comprise searching for the transforma-

tion that does not change an orientation of the transformed tracking map with respect to the gravity direction.

[0021] According to one embodiment, searching for a transformation may comprise determining whether applying the transformation to the tracking map produces the transformed tracking map that is rotated with respect to the gravity direction by more than a threshold amount; in response to determining that applying the transformation to the tracking map does not produce the transformed tracking map that is rotated with respect to the gravity direction by more than the threshold amount, applying the transformation to the tracking map to generate the transformed tracking map; and in response to determining that applying the transformation to the tracking map produces the transformed tracking map that is rotated with respect to the gravity direction by more than the threshold amount, discarding the transformation.

[0022] According to one embodiment, the method may further comprise identifying a set of environment maps from the plurality of environment maps to be merged with the tracking map; and for each environment map in the set of environment maps: determining whether to merge the environment map with the tracking map; and merging the environment map with the transformed tracking map based on determining that a gravity direction of the environment map aligns with the gravity direction of the transformed tracking map.

[0023] According to one embodiment, the method may further comprise, for each environment map in the set of environment maps: refraining from merging the environment map with the transformed tracking map based on determining that the gravity direction of the environment map does not align with the gravity direction of the tracking map.

[0024] According to one embodiment, identifying the set of environment maps may further comprise determining area identifiers associated with the tracking map; and identifying the set of environment maps based, at least in part on, on the area identifiers associated with the tracking map.

[0025] According to one embodiment, identifying the set of environment maps may further comprise filtering the set of environment maps based on similarity of one or more metrics associated with the tracking map and the environment maps in the set of environment maps.

[0026] According to one embodiment, a cloud computing environment for an augmented reality system is configured for communication with a plurality of user devices comprising sensors. The cloud computing environment for an augmented reality system may comprise a map database storing a plurality of environment maps constructed from data supplied by the plurality of user devices; and non-transitory computer storage media storing computer-executable instructions that, when executed by at least one processor in the cloud computing environment, perform a method. The method may comprise receiving a tracking map from a user device, wherein the tracking map is aligned with respect to a gravity direction; updating the map database based on the received tracking map, wherein updating the map database comprises, for each environment map in a set of environment maps: determining a transformation between the tracking map and the environment map; determining whether to merge the environment map with the tracking map, wherein determining whether to merge comprises determining whether applying the transformation to the tracking map

produced a transformed tracking map that aligns with respect to the gravity direction; and merging the environment map with the tracking map based on determining that the transformed tracking map aligns with the gravity direction.

[0027] According to one embodiment, determining the transformation may comprise, for corresponding features in the tracking map and the environment map, selecting as the determined transformation, the transformation that aligns the corresponding features with a metric of error below a threshold.

[0028] According to one embodiment, the method may further comprise determining the corresponding features based on similarity of identifiers assigned to the features.

[0029] According to one embodiment, selecting as the determined transformation may further comprise selecting the transformation that when applied produces the transformed tracking map that aligns with respect to the gravity direction.

[0030] According to one embodiment, determining the transformation may comprise applying a plurality of candidate transformations to the tracking map and selecting as the determined transformation a candidate transformation of the plurality of candidate transformations.

[0031] According to one embodiment, selecting as the determined transformation may further comprise selecting as the determined transformation the candidate transformation that when applied produces the transformed tracking map that aligns with respect to the gravity direction.

[0032] According to one embodiment, determining whether applying the transformation to the tracking map produces the transformed tracking map that aligns with respect to the gravity direction may comprise: determining whether applying the transformation to the tracking map produces the transformed tracking map that is rotated with respect to the gravity direction by more than a threshold amount; in response to determining that applying the transformation to the tracking map does not produce the transformed tracking map that is rotated with respect to the gravity direction by more than the threshold amount, selecting the transformation as the determined transformation; and in response to determining that applying the transformation to the tracking map produces the transformed tracking map that is rotated with respect to the gravity direction by more than the threshold amount, discarding the transformation.

[0033] The foregoing summary is provided by way of illustration and is not intended to be limiting.

BRIEF DESCRIPTION OF THE DRAWINGS

[0034] The accompanying drawings are not intended to be drawn to scale. In the drawings, each identical or nearly identical component that is illustrated in various figures is represented by a like numeral. For purposes of clarity, not every component may be labeled in every drawing. In the drawings:

[0035] FIG. 1 is a sketch illustrating an example of a simplified augmented reality (AR) scene, according to some embodiments;

[0036] FIG. 2 is a sketch of an exemplary simplified AR scene, showing exemplary use cases of an XR system, according to some embodiments;

[0037] FIG. 3 is a schematic diagram illustrating data flow for a single user in an AR system configured to provide an

experience to the user of AR content interacting with a physical world, according to some embodiments;

[0038] FIG. 4 is a schematic diagram illustrating an exemplary AR display system, displaying virtual content for a single user, according to some embodiments;

[0039] FIG. 5A is a schematic diagram illustrating a user wearing an AR display system rendering AR content as the user moves through a physical world environment, according to some embodiments;

[0040] FIG. 5B is a schematic diagram illustrating a viewing optics assembly and attendant components, according to some embodiments;

[0041] FIG. 6A is a schematic diagram illustrating an AR system using a world reconstruction system, according to some embodiments;

[0042] FIG. 6B is a schematic diagram illustrating components of an AR system that maintain a model of a passable world, according to some embodiments;

[0043] FIG. 7 is a schematic illustration of a tracking map formed by a device traversing a path through a physical world;

[0044] FIG. 8 is a schematic diagram illustrating a user of a cross reality (XR) system, perceiving virtual content, according to some embodiments;

[0045] FIG. 9 is a block diagram of components of a first XR device of the XR system of FIG. 8 that transform between coordinate systems, according to some embodiments;

[0046] FIG. 10 is a schematic diagram illustrating an exemplary transformation of origin coordinate frames into destination coordinate frames in order to correctly render local XR content, according to some embodiments;

[0047] FIG. 11 is a top plan view illustrating pupil-based coordinate frames, according to some embodiments;

[0048] FIG. 12 is a top plan view illustrating a camera coordinate frame that includes all pupil positions, according to some embodiments;

[0049] FIG. 13 is a schematic diagram of the display system of FIG. 9, according to some embodiments;

[0050] FIG. 14 is a block diagram illustrating the creation of a persistent coordinate frame (PCF) and the attachment of XR content to the PCF, according to some embodiments;

[0051] FIG. 15 is a flow chart illustrating a method of establishing and using a PCF, according to some embodiments;

[0052] FIG. 16 is a block diagram of the XR system of FIG. 8, including a second XR device, according to some embodiments;

[0053] FIG. 17 is a schematic diagram illustrating a room and key frames that are established for various areas in the room, according to some embodiments;

[0054] FIG. 18 is a schematic diagram illustrating the establishment of persistent poses based on the key frames, according to some embodiments;

[0055] FIG. 19 is a schematic diagram illustrating the establishment of a persistent coordinate frame (PCF) based on the persistent poses, according to some embodiments;

[0056] FIGS. 20A to 20C are schematic diagrams illustrating an example of creating PCFs, according to some embodiments;

[0057] FIG. 21 is a block diagram illustrating a system for generating global descriptors for individual images and/or maps, according to some embodiments;

- [0058] FIG. 22 is a flow chart illustrating a method of computing an image descriptor, according to some embodiments;
- [0059] FIG. 23 is a flow chart illustrating a method of localization using image descriptors, according to some embodiments;
- [0060] FIG. 24 is a flow chart illustrating a method of training a neural network, according to some embodiments;
- [0061] FIG. 25 is a block diagram illustrating a method of training a neural network, according to some embodiments;
- [0062] FIG. 26 is a schematic diagram illustrating an AR system configured to rank and merge a plurality of environment maps, according to some embodiments;
- [0063] FIG. 27 is a simplified block diagram illustrating a plurality of canonical maps stored on a remote storage medium, according to some embodiments;
- [0064] FIG. 28 is a schematic diagram illustrating a method of selecting canonical maps to, for example, localize a new tracking map in one or more canonical maps and/or obtain PCFs from the canonical maps, according to some embodiments;
- [0065] FIG. 29 is flow chart illustrating a method of selecting a plurality of ranked environment maps, according to some embodiments;
- [0066] FIG. 30 is a schematic diagram illustrating an exemplary map rank portion of the AR system of FIG. 26, according to some embodiments;
- [0067] FIG. 31A is a schematic diagram illustrating an example of area attributes of a tracking map (TM) and environment maps in a database, according to some embodiments;
- [0068] FIG. 31B is a schematic diagram illustrating an example of determining a geographic location of a tracking map (TM) for geolocation filtering of FIG. 29, according to some embodiments;
- [0069] FIG. 32 is a schematic diagram illustrating an example of geolocation filtering of FIG. 29, according to some embodiments;
- [0070] FIG. 33 is a schematic diagram illustrating an example of Wi-Fi BSSID filtering of FIG. 29, according to some embodiments;
- [0071] FIG. 34 is a schematic diagram illustrating an example of use of localization of FIG. 29, according to some embodiments;
- [0072] FIGS. 35 and 36 are block diagrams of an XR system configured to rank and merge a plurality of environment maps, according to some embodiments.
- [0073] FIG. 37 is a block diagram illustrating a method of creating environment maps of a physical world, in a canonical form, according to some embodiments;
- [0074] FIGS. 38A and 38B are schematic diagrams illustrating an environment map created in a canonical form by updating the tracking map of FIG. 7 with a new tracking map, according to some embodiments.
- [0075] FIGS. 39A to 39F are schematic diagrams illustrating an example of merging maps, according to some embodiments;
- [0076] FIG. 40 is a two-dimensional representation of a three-dimensional first local tracking map (Map 1), which may be generated by the first XR device of FIG. 9, according to some embodiments;
- [0077] FIG. 41 is a block diagram illustrating uploading Map 1 from the first XR device to the server of FIG. 9, according to some embodiments;
- [0078] FIG. 42 is a schematic diagram illustrating the XR system of FIG. 16, showing the second user has initiated a second session using a second XR device of the XR system after the first user has terminated a first session, according to some embodiments;
- [0079] FIG. 43A is a block diagram illustrating a new session for the second XR device of FIG. 42, according to some embodiments;
- [0080] FIG. 43B is a block diagram illustrating the creation of a tracking map for the second XR device of FIG. 42, according to some embodiments;
- [0081] FIG. 43C is a block diagram illustrating downloading a canonical map from the server to the second XR device of FIG. 42, according to some embodiments;
- [0082] FIG. 44 is a schematic diagram illustrating a localization attempt to localize to a canonical map a second tracking map (Map 2), which may be generated by the second XR device of FIG. 42, according to some embodiments;
- [0083] FIG. 45 is a schematic diagram illustrating a localization attempt to localize to a canonical map the second tracking map (Map 2) of FIG. 44, which may be further developed and with XR content associated with PCFs of Map 2, according to some embodiments;
- [0084] FIGS. 46A-46B are a schematic diagram illustrating a successful localization of Map 2 of FIG. 45 to the canonical map, according to some embodiments;
- [0085] FIG. 47 is a schematic diagram illustrating a canonical map generated by including one or more PCFs from the canonical map of FIG. 46A into Map 2 of FIG. 45, according to some embodiments;
- [0086] FIG. 48 is a schematic diagram illustrating the canonical map of FIG. 47 with further expansion of Map 2 on the second XR device, according to some embodiments;
- [0087] FIG. 49 is a block diagram illustrating uploading Map 2 from the second XR device to the server, according to some embodiments;
- [0088] FIG. 50 is a block diagram illustrating merging Map 2 with the canonical map, according to some embodiments;
- [0089] FIG. 51 is a block diagram illustrating transmission of a new canonical map from the server to the first and second XR devices, according to some embodiments;
- [0090] FIG. 52 is a block diagram illustrating a two-dimensional representation of Map 2 and a head coordinate frame of the second XR device that is referenced to Map 2, according to some embodiments;
- [0091] FIG. 53 is a block diagram illustrating, in two-dimensions, adjustment of the head coordinate frame which can occur in six degrees of freedom, according to some embodiments;
- [0092] FIG. 54 is a block diagram illustrating a canonical map on the second XR device wherein sound is localized relative to PCFs of Map 2, according to some embodiments;
- [0093] FIGS. 55 and 56 are a perspective view and a block diagram illustrating use of the XR system when the first user has terminated a first session and the first user has initiated a second session using the XR system, according to some embodiments;
- [0094] FIGS. 57 and 58 are a perspective view and a block diagram illustrating use of the XR system when three users are simultaneously using the XR system in the same session, according to some embodiments;

[0095] FIG. 59 is a flow chart illustrating a method of recovering and resetting a headpose, according to some embodiments;

[0096] FIG. 60 is a block diagram of a machine in the form of a computer that can find application in the present invention system, according to some embodiments;

[0097] FIG. 61 is a schematic diagram of an example XR system in which any of multiple devices may access a localization service, according to some embodiments;

[0098] FIG. 62 is an example process flow for operation of a portable device as part of an XR system that provides cloud-based localization, according to some embodiments;

[0099] FIGS. 63A, B, and C are an example process flow for cloud-based localization, according to some embodiments;

[0100] FIG. 64 is a block diagram of an XR system that provides large scale localization, according to some embodiments;

[0101] FIG. 65 is a schematic diagram illustrating information about a physical world being processed by the XR system of FIG. 64, according to some embodiments;

[0102] FIG. 66 is a block diagram of a subsystem of the XR system of FIG. 64 including the matched correspondences quality predication component and a pose estimation component, according to some embodiments;

[0103] FIG. 67 is a flow chart illustrating a method of generating dataset for training the subsystem of FIG. 66, according to some embodiments;

[0104] FIG. 68 is a flow chart illustrating a method of training the subsystem of FIG. 66, according to some embodiments; and

[0105] FIG. 69 illustrates a gravity-preserving map merge process, according to some embodiments.

DETAILED DESCRIPTION

[0106] Described herein are methods and apparatus for providing XR scenes. To provide realistic XR experiences to multiple users, an XR system must know the users' location within the physical world in order to correctly correlate locations of virtual objects in relation to real objects. The inventors have recognized and appreciated methods and apparatus that localize XR devices in large and very large scale environments (e.g., a neighborhood, a city, a country, the globe) with reduced time and improved accuracy.

[0107] An XR system may build an environment map of a scene, which may be created from image and/or depth information collected with sensors that are part of XR devices worn by users of the XR system. Each XR device may develop a local map of its physical environment by integrating information from one or more images collected as the device operates. In some embodiments, the coordinate system of that map is tied to the orientation of the device when the device initiated scanning the physical world. That orientation may change from session to session as a user interacts with the XR system, whether different sessions are associated with different users, each with their own wearable device with sensors that scan the environment, or the same user who uses the same device at different times.

[0108] The XR system may implement one or more techniques so as to enable operation based on persistent spatial information. The techniques, for example, may provide XR scenes for a more computationally efficient and immersive experience for a single or multiple users by enabling persistent spatial information to be created, stored, and

retrieved by any of multiple users of an XR system. Persistent spatial information may also enable quickly recovering and resetting headposes on each of one or more XR devices in a computationally efficient way.

[0109] The persistent spatial information may be represented by a persistent map. The persistent map may be stored in a remote storage medium (e.g., a cloud). For example, the wearable device worn by a user, after being turned on, may retrieve from persistent storage, such as from cloud storage, an appropriate map that was previously created and stored. That previously stored map may have been based on data about the environment collected with sensors on the user's wearable device during prior sessions. Retrieving a stored map may enable use of the wearable device without completing a scan of the physical world with the sensors on the wearable device. Alternatively or additionally, the system/device, upon entering a new region of the physical world, may similarly retrieve an appropriate stored map.

[0110] The stored map may be represented in a canonical form to which a local frame of reference on each XR device may be related. In a multidevice XR system, the stored map accessed by one device may have been created and stored by another device and/or may have been constructed by aggregating data about the physical world collected by sensors on multiple wearable devices that were previously present in at least a portion of the physical world represented by the stored map.

[0111] In some embodiments, persistent spatial information may be represented in a way that may be readily shared among users and among the distributed components, including applications. Canonical maps may provide information about the physical world, for example, as persistent coordinate frames (PCFs). A PCF may be defined based on a set of features recognized in the physical world. The features may be selected such that they are likely to be the same from user session to user session of the XR system. PCFs may exist sparsely, providing less than all of the available information about the physical world, such that they may be efficiently processed and transferred. Techniques for processing persistent spatial information may include creating dynamic maps based on the local coordinate systems of one or more devices across one or more sessions. These maps may be sparse maps, representing the physical world based on a subset of the feature points detected in images used in forming the maps. The persistent coordinate frames (PCF) may be generated from the sparse maps, and may be exposed to XR applications via, for example, an application programming interface (API). These capabilities may be supported by techniques for forming the canonical maps by merging multiple maps created by one or more XR devices.

[0112] The relationship between the canonical map and a local map for each device may be determined through a localization process. That localization process may be performed on each XR device based on a set of canonical maps selected and sent to the device. Alternatively or additionally, a localization service may be provided on remote processors, such as might be implemented in the cloud.

[0113] Sharing data about the physical world among multiple devices may enable shared user experiences of virtual content. Two XR devices that have access to the same stored map, for example, may both localize with respect to the stored map. Once localized, a user device may render virtual content that has a location specified by reference to the stored map by translating that location to a frame of refer-

ence maintained by the user device. The user device may use this local frame of reference to control the display of the user device to render the virtual content in the specified location.

[0114] To support these and other functions, the XR system may include components that, based on data about the physical world collected with sensors on user devices, develop, maintain, and use persistent spatial information, including one or more stored maps. These components may be distributed across the XR system, with some operating, for example, on a head mounted portion of a user device. Other components may operate on a computer, associated with the user coupled to the head mounted portion over a local or personal area network. Yet others may operate at a remote location, such as at one or more servers accessible over a wide area network.

[0115] These components, for example, may include components that can identify from information about the physical world collected by one or more user devices information that is of sufficient quality to be stored as or in a persistent map. An example of such a component, described in greater detail below, is a map merge component. Such a component, for example, may receive inputs from a user device and determine the suitability of parts of the inputs to be used to update a persistent map, which may be in canonical form. A map merge component, for example, may also promote a local map from a user device that is not merged with a persistent map to be a separate persistent map.

[0116] As another example, these components may include components that may aid in selecting an appropriate set of one or more persistent maps that likely represent the same region of the physical world as is represented by location information provided by a user device. An example of such components, described in greater detail below are map rank and map select components. Such components, for example, may receive inputs from a user device and identify one or more persistent maps that are likely to represent the region of the physical world in which that device is operating. A map rank component, for example, may aid in selecting a persistent map to be used by that local device as it renders virtual content, gathers data about the environment, or performs other actions. A map rank component, alternatively or additionally, may aid in identifying persistent maps to be updated as additional information about the physical world is collected by one or more user devices.

[0117] Yet other components may determine transformations that transform information captured or described in relation to one reference frame into another reference frame. For example, sensors may be attached to a head mounted display such that the data read from those sensors indicates locations of objects in the physical world with respect to the headpose of the wearer. One or more transformations may be applied to relate that location information to the coordinate frame associated with a persistent environment map. Similarly, data indicating where a virtual object is to be rendered when expressed in a coordinate frame of a persistent environment map may be put through one or more transformations to be in a frame of reference of the display on the user's head. As described in greater detail below, there may be multiple such transformations. These transformations may be partitioned across the components of an XR system such that they may be efficiently updated and or applied in a distributed system.

[0118] In some embodiments, the persistent maps may be constructed from information collected by multiple user

devices. The XR devices may each capture local spatial information and construct separate tracking maps with information collected by sensors of each of the XR devices at various locations and times. Each tracking map may include points, each of which may be associated with a feature of a real object that may include multiple features. In addition to potentially supplying input to create and maintain persistent maps, the tracking maps may be used to track users' motions in a scene, enabling an XR system to estimate respective users' headposes relative to the frame of reference established by the tracking map on that user's device.

[0119] This co-dependence between the creation of a map and the estimation of headpose constitutes significant challenges. Substantial processing may be required to create the map and estimate headposes simultaneously. The processing must be accomplished quickly as objects move in the scene (e.g., moving a cup on a table) and as users move in the scene because latency makes XR experiences less realistic for users. On the other hand, an XR device can provide limited computational resources because an XR device should be lightweight for a user to wear comfortably. Lack of computational resources cannot be compensated for with more sensors, as adding sensors would also undesirably add weight. Further, either more sensors or more computational resources leads to heat, which may cause deformation of an XR device.

[0120] The XR system may be configured to create, share, and use persistent spatial information with low usage of computational resources and/or low latency to provide a more immersive user experience. Some such techniques may enable efficient comparison of spatial information. Such comparisons may arise, for example, as part of localization in which a set of features from a local device is matched to a set of features in a canonical map.

[0121] Similarly, in a map merge process, attempts may be made to match one or more sets of features in a tracking map from a device to corresponding features in a canonical map and determine a transformation between the sets of corresponding features that provides a suitably low error between the positions of transformed features in a first set of features derived from the tracking map and a second set of features derived from the canonical map. Subsequent processing to incorporate the tracking map into a set of canonical maps may be based on the results of that comparison. For example, determining a transformation with suitably low error may indicate that a region represented by the second set of features derived from the canonical map correspond to a same region represented by the first set of features derived from the tracking map and that the two maps can be merged.

[0122] The inventors have recognized that errors may, nonetheless, be introduced in the merging process even when there is low error in the alignment of the sets of features. The inventors have further recognized and appreciated that such errors may be detected by a transformation that changes the orientation of the tracking map with respect to a gravity direction, and that merging such a tracking map with the canonical map may result in a skewed merged map.

[0123] By suppressing the merging of transformed tracking maps that have an orientation changed with respect to gravity, the canonical map may start with, and retain an alignment with respect to gravity. Merging errors can be reduced by ensuring that a gravity direction of a transformed tracking map (i.e., after applying the determined transfor-

mation) aligns with a gravity direction of the canonical map with which the tracking map is to be merged.

[0124] Techniques as described herein may be used together or separately with many types of devices and for many types of scenes, including wearable or portable devices with limited computational resources that provide an augmented or mixed reality scene. In some embodiments, the techniques may be implemented by one or more services that form a portion of an XR system.

AR System Overview

[0125] FIGS. 1 and 2 illustrate scenes with virtual content displayed in conjunction with a portion of the physical world. For purposes of illustration, an AR system is used as an example of an XR system. FIGS. 3-6B illustrate an exemplary AR system, including one or more processors, memory, sensors and user interfaces that may operate according to the techniques described herein.

[0126] Referring to FIG. 1, an outdoor AR scene 354 is depicted in which a user of an AR technology sees a physical world park-like setting 356, featuring people, trees, buildings in the background, and a concrete platform 358. In addition to these items, the user of the AR technology also perceives that they “see” a robot statue 357 standing upon the physical world concrete platform 358, and a cartoon-like avatar character 352 flying by which seems to be a personification of a bumble bee, even though these elements (e.g., the avatar character 352, and the robot statue 357) do not exist in the physical world. Due to the extreme complexity of the human visual perception and nervous system, it is challenging to produce an AR technology that facilitates a comfortable, natural-feeling, rich presentation of virtual image elements amongst other virtual or physical world imagery elements.

[0127] Such an AR scene may be achieved with a system that builds maps of the physical world based on tracking information, enable users to place AR content in the physical world, determine locations in the maps of the physical world where AR content are placed, preserve the AR scenes such that the placed AR content can be reloaded to display in the physical world during, for example, a different AR experience session, and enable multiple users to share an AR experience. The system may build and update a digital representation of the physical world surfaces around the user. This representation may be used to render virtual content so as to appear fully or partially occluded by physical objects between the user and the rendered location of the virtual content, to place virtual objects, in physics based interactions, and for virtual character path planning and navigation, or for other operations in which information about the physical world is used.

[0128] FIG. 2 depicts another example of an indoor AR scene 400, showing exemplary use cases of an XR system, according to some embodiments. The exemplary scene 400 is a living room having walls, a bookshelf on one side of a wall, a floor lamp at a corner of the room, a floor, a sofa, and coffee table on the floor. In addition to these physical items, the user of the AR technology also perceives virtual objects such as images on the wall behind the sofa, birds flying through the door, a deer peeking out from the book shelf, and a decoration in the form of a windmill placed on the coffee table.

[0129] For the images on the wall, the AR technology requires information about not only surfaces of the wall but

also objects and surfaces in the room such as lamp shape, which are occluding the images to render the virtual objects correctly. For the flying birds, the AR technology requires information about all the objects and surfaces around the room for rendering the birds with realistic physics to avoid the objects and surfaces or bounce off them if the birds collide. For the deer, the AR technology requires information about the surfaces such as the floor or coffee table to compute where to place the deer. For the windmill, the system may identify that is an object separate from the table and may determine that it is movable, whereas corners of shelves or corners of the wall may be determined to be stationary. Such a distinction may be used in determinations as to which portions of the scene are used or updated in each of various operations.

[0130] The virtual objects may be placed in a previous AR experience session. When new AR experience sessions start in the living room, the AR technology requires the virtual objects being accurately displayed at the locations previously placed and realistically visible from different viewpoints. For example, the windmill should be displayed as standing on the books rather than drifting above the table at a different location without the books. Such drifting may happen if the locations of the users of the new AR experience sessions are not accurately localized in the living room. As another example, if a user is viewing the windmill from a viewpoint different from the viewpoint when the windmill was placed, the AR technology requires corresponding sides of the windmill being displayed.

[0131] A scene may be presented to the user via a system that includes multiple components, including a user interface that can stimulate one or more user senses, such as sight, sound, and/or touch. In addition, the system may include one or more sensors that may measure parameters of the physical portions of the scene, including position and/or motion of the user within the physical portions of the scene. Further, the system may include one or more computing devices, with associated computer hardware, such as memory. These components may be integrated into a single device or may be distributed across multiple interconnected devices. In some embodiments, some or all of these components may be integrated into a wearable device.

[0132] FIG. 3 depicts an AR system 502 configured to provide an experience of AR contents interacting with a physical world 506, according to some embodiments. The AR system 502 may include a display 508. In the illustrated embodiment, the display 508 may be worn by the user as part of a headset such that a user may wear the display over their eyes like a pair of goggles or glasses. At least a portion of the display may be transparent such that a user may observe a see-through reality 510. The see-through reality 510 may correspond to portions of the physical world 506 that are within a present viewpoint of the AR system 502, which may correspond to the viewpoint of the user in the case that the user is wearing a headset incorporating both the display and sensors of the AR system to acquire information about the physical world.

[0133] AR contents may also be presented on the display 508, overlaid on the see-through reality 510. To provide accurate interactions between AR contents and the see-through reality 510 on the display 508, the AR system 502 may include sensors 522 configured to capture information about the physical world 506.

[0134] The sensors 522 may include one or more depth sensors that output depth maps 512. Each depth map 512 may have multiple pixels, each of which may represent a distance to a surface in the physical world 506 in a particular direction relative to the depth sensor. Raw depth data may come from a depth sensor to create a depth map. Such depth maps may be updated as fast as the depth sensor can form a new image, which may be hundreds or thousands of times per second. However, that data may be noisy and incomplete, and have holes shown as black pixels on the illustrated depth map.

[0135] The system may include other sensors, such as image sensors. The image sensors may acquire monocular or stereoscopic information that may be processed to represent the physical world in other ways. For example, the images may be processed in world reconstruction component 516 to create a mesh, representing connected portions of objects in the physical world. Metadata about such objects, including for example, color and surface texture, may similarly be acquired with the sensors and stored as part of the world reconstruction.

[0136] The system may also acquire information about the headpose (or “pose”) of the user with respect to the physical world. In some embodiments, a headpose tracking component of the system may be used to compute headposes in real time. The headpose tracking component may represent a headpose of a user in a coordinate frame with six degrees of freedom including, for example, translation in three perpendicular axes (e.g., forward/backward, up/down, left/right) and rotation about the three perpendicular axes (e.g., pitch, yaw, and roll). In some embodiments, sensors 522 may include inertial measurement units that may be used to compute and/or determine a headpose 514. A headpose 514 for a depth map may indicate a present viewpoint of a sensor capturing the depth map with six degrees of freedom, for example, but the headpose 514 may be used for other purposes, such as to relate image information to a particular portion of the physical world or to relate the position of the display worn on the user’s head to the physical world.

[0137] In some embodiments, the headpose information may be derived in other ways than from an IMU, such as from analyzing objects in an image. For example, the headpose tracking component may compute relative position and orientation of an AR device to physical objects based on visual information captured by cameras and inertial information captured by IMUs. The headpose tracking component may then compute a headpose of the AR device by, for example, comparing the computed relative position and orientation of the AR device to the physical objects with features of the physical objects. In some embodiments, that comparison may be made by identifying features in images captured with one or more of the sensors 522 that are stable over time such that changes of the position of these features in images captured over time can be associated with a change in headpose of the user.

[0138] The inventors have realized and appreciated techniques for operating XR systems to provide XR scenes for a more immersive user experience such as estimating headpose at a frequency of 1 kHz, with low usage of computational resources in connection with an XR device, that may be configured with, for example, four video graphic array (VGA) cameras operating at 30 Hz, one inertial measurement unit (IMU) operating at 1 kHz, compute power of a single advanced RISC machine (ARM) core, memory less

than 1 GB, and network bandwidth less than 100 Mbp. These techniques relate to reducing processing required to generate and maintain maps and estimate headpose as well as to providing and consuming data with low computational overhead. The XR system may calculate its pose based on the matched visual features. U.S. patent application Ser. No. 16/221,065 describes hybrid tracking and is hereby incorporated herein by reference in its entirety.

[0139] In some embodiments, the AR device may construct a map from the feature points recognized in successive images in a series of image frames captured as a user moves throughout the physical world with the AR device. Though each image frame may be taken from a different pose as the user moves, the system may adjust the orientation of the features of each successive image frame to match the orientation of the initial image frame by matching features of the successive image frames to previously captured image frames. Translations of the successive image frames so that points representing the same features will match corresponding feature points from previously collected image frames, can be used to align each successive image frame to match the orientation of previously processed image frames. The frames in the resulting map may have a common orientation established when the first image frame was added to the map. This map, with sets of feature points in a common frame of reference, may be used to determine the user’s pose within the physical world by matching features from current image frames to the map. In some embodiments, this map may be called a tracking map.

[0140] In addition to enabling tracking of the user’s pose within the environment, this map may enable other components of the system, such as world reconstruction component 516, to determine the location of physical objects with respect to the user. The world reconstruction component 516 may receive the depth maps 512 and headposes 514, and any other data from the sensors, and integrate that data into a reconstruction 518. The reconstruction 518 may be more complete and less noisy than the sensor data. The world reconstruction component 516 may update the reconstruction 518 using spatial and temporal averaging of the sensor data from multiple viewpoints over time.

[0141] The reconstruction 518 may include representations of the physical world in one or more data formats including, for example, voxels, meshes, planes, etc. The different formats may represent alternative representations of the same portions of the physical world or may represent different portions of the physical world. In the illustrated example, on the left side of the reconstruction 518, portions of the physical world are presented as a global surface; on the right side of the reconstruction 518, portions of the physical world are presented as meshes.

[0142] In some embodiments, the map maintained by headpose component 514 may be sparse relative to other maps that might be maintained of the physical world. Rather than providing information about locations, and possibly other characteristics, of surfaces, the sparse map may indicate locations of interest points and/or structures, such as corners or edges. In some embodiments, the map may include image frames as captured by the sensors 522. These frames may be reduced to features, which may represent the interest points and/or structures. In conjunction with each frame, information about a pose of a user from which the frame was acquired may also be stored as part of the map. In some embodiments, every image acquired by the sensor

may or may not be stored. In some embodiments, the system may process images as they are collected by sensors and select subsets of the image frames for further computation. The selection may be based on one or more criteria that limits the addition of information yet ensures that the map contains useful information. The system may add a new image frame to the map, for example, based on overlap with a prior image frame already added to the map or based on the image frame containing a sufficient number of features determined as likely to represent stationary objects. In some embodiments, the selected image frames, or groups of features from selected image frames may serve as key frames for the map, which are used to provide spatial information.

[0143] In some embodiments, the amount of data that is processed when constructing maps may be reduced, such as by constructing sparse maps with a collection of mapped points and keyframes and/or dividing the maps into blocks to enable updates by blocks. A mapped point may be associated with a point of interest in the environment. A keyframe may include selected information from camera-captured data. U.S. patent application Ser. No. 16/520,582 describes determining and/or evaluating localization maps and is hereby incorporated herein by reference in its entirety.

[0144] The AR system 502 may integrate sensor data over time from multiple viewpoints of a physical world. The poses of the sensors (e.g., position and orientation) may be tracked as a device including the sensors is moved. As the sensor's frame pose is known and how it relates to the other poses, each of these multiple viewpoints of the physical world may be fused together into a single, combined reconstruction of the physical world, which may serve as an abstract layer for the map and provide spatial information. The reconstruction may be more complete and less noisy than the original sensor data by using spatial and temporal averaging (i.e. averaging data from multiple viewpoints over time), or any other suitable method.

[0145] In the illustrated embodiment in FIG. 3, a map represents the portion of the physical world in which a user of a single, wearable device is present. In that scenario, headpose associated with frames in the map may be represented as a local headpose, indicating orientation relative to an initial orientation for a single device at the start of a session. For example, the headpose may be tracked relative to an initial headpose when the device was turned on or otherwise operated to scan an environment to build a representation of that environment.

[0146] In combination with content characterizing that portion of the physical world, the map may include metadata. The metadata, for example, may indicate time of capture of the sensor information used to form the map. Metadata alternatively or additionally may indicate location of the sensors at the time of capture of information used to form the map. Location may be expressed directly, such as with information from a GPS chip, or indirectly, such as with a wireless (e.g. Wi-Fi) signature indicating strength of signals received from one or more wireless access points while the sensor data was being collected and/or with identifiers, such as BSSID's, of wireless access points to which the user device connected while the sensor data was collected.

[0147] The reconstruction 518 may be used for AR functions, such as producing a surface representation of the physical world for occlusion processing or physics-based

processing. This surface representation may change as the user moves or objects in the physical world change. Aspects of the reconstruction 518 may be used, for example, by a component 520 that produces a changing global surface representation in world coordinates, which may be used by other components.

[0148] The AR content may be generated based on this information, such as by AR applications 504. An AR application 504 may be a game program, for example, that performs one or more functions based on information about the physical world, such as visual occlusion, physics-based interactions, and environment reasoning. It may perform these functions by querying data in different formats from the reconstruction 518 produced by the world reconstruction component 516. In some embodiments, component 520 may be configured to output updates when a representation in a region of interest of the physical world changes. That region of interest, for example, may be set to approximate a portion of the physical world in the vicinity of the user of the system, such as the portion within the view field of the user, or is projected (predicted/determined) to come within the view field of the user.

[0149] The AR applications 504 may use this information to generate and update the AR contents. The virtual portion of the AR contents may be presented on the display 508 in combination with the see-through reality 510, creating a realistic user experience.

[0150] In some embodiments, an AR experience may be provided to a user through an XR device, which may be a wearable display device, which may be part of a system that may include remote processing and/or remote data storage and/or, in some embodiments, other wearable display devices worn by other users. FIG. 4 illustrates an example of system 580 (hereinafter referred to as "system 580") including a single wearable device for simplicity of illustration. The system 580 includes a head mounted display device 562 (hereinafter referred to as "display device 562"), and various mechanical and electronic modules and systems to support the functioning of the display device 562. The display device 562 may be coupled to a frame 564, which is wearable by a display system user or viewer 560 (hereinafter referred to as "user 560") and configured to position the display device 562 in front of the eyes of the user 560. According to various embodiments, the display device 562 may be a sequential display. The display device 562 may be monocular or binocular. In some embodiments, the display device 562 may be an example of the display 508 in FIG. 3.

[0151] In some embodiments, a speaker 566 is coupled to the frame 564 and positioned proximate an ear canal of the user 560. In some embodiments, another speaker, not shown, is positioned adjacent another ear canal of the user 560 to provide for stereo/shapeable sound control. The display device 562 is operatively coupled, such as by a wired lead or wireless connectivity 568, to a local data processing module 570 which may be mounted in a variety of configurations, such as fixedly attached to the frame 564, fixedly attached to a helmet or hat worn by the user 560, embedded in headphones, or otherwise removably attached to the user 560 (e.g., in a backpack-style configuration, in a belt-coupling style configuration).

[0152] The local data processing module 570 may include a processor, as well as digital memory, such as non-volatile memory (e.g., flash memory), both of which may be utilized to assist in the processing, caching, and storage of data. The

data include data a) captured from sensors (which may be, e.g., operatively coupled to the frame 564) or otherwise attached to the user 560, such as image capture devices (such as cameras), microphones, inertial measurement units, accelerometers, compasses, GPS units, radio devices, and/or gyros; and/or b) acquired and/or processed using remote processing module 572 and/or remote data repository 574, possibly for passage to the display device 562 after such processing or retrieval.

[0153] In some embodiments, the wearable device may communicate with remote components. The local data processing module 570 may be operatively coupled by communication links 576, 578, such as via a wired or wireless communication links, to the remote processing module 572 and remote data repository 574, respectively, such that these remote modules 572, 574 are operatively coupled to each other and available as resources to the local data processing module 570. In further embodiments, in addition or as alternative to remote data repository 574, the wearable device can access cloud based remote data repositories, and/or services. In some embodiments, the headpose tracking component described above may be at least partially implemented in the local data processing module 570. In some embodiments, the world reconstruction component 516 in FIG. 3 may be at least partially implemented in the local data processing module 570. For example, the local data processing module 570 may be configured to execute computer executable instructions to generate the map and/or the physical world representations based at least in part on at least a portion of the data.

[0154] In some embodiments, processing may be distributed across local and remote processors. For example, local processing may be used to construct a map on a user device (e.g. tracking map) based on sensor data collected with sensors on that user's device. Such a map may be used by applications on that user's device. Additionally, previously created maps (e.g., canonical maps) may be stored in remote data repository 574. Where a suitable stored or persistent map is available, it may be used instead of or in addition to the tracking map created locally on the device. In some embodiments, a tracking map may be localized to the stored map, such that a correspondence is established between a tracking map, which might be oriented relative to a position of the wearable device at the time a user turned the system on, and the canonical map, which may be oriented relative to one or more persistent features. In some embodiments, the persistent map might be loaded on the user device to allow the user device to render virtual content without a delay associated with scanning a location to build a tracking map of the user's full environment from sensor data acquired during the scan. In some embodiments, the user device may access a remote persistent map (e.g., stored on a cloud) without the need to download the persistent map on the user device.

[0155] In some embodiments, spatial information may be communicated from the wearable device to remote services, such as a cloud service that is configured to localize a device to stored maps maintained on the cloud service. According to one embodiment, the localization processing can take place in the cloud matching the device location to existing maps, such as canonical maps, and return transforms that link virtual content to the wearable device location. In such embodiments, the system can avoid communicating maps from remote resources to the wearable device. Other

embodiments can be configured for both device-based and cloud-based localization, for example, to enable functionality where network connectivity is not available or a user opts not to enable cloud-based localization.

[0156] Alternatively or additionally, the tracking map may be merged with previously stored maps to extend or improve the quality of those maps. The processing to determine whether a suitable previously created environment map is available and/or to merge a tracking map with one or more stored environment maps may be done in local data processing module 570 or remote processing module 572.

[0157] In some embodiments, the local data processing module 570 may include one or more processors (e.g., a graphics processing unit (GPU)) configured to analyze and process data and/or image information. In some embodiments, the local data processing module 570 may include a single processor (e.g., a single-core or multi-core ARM processor), which would limit the local data processing module 570's compute budget but enable a more miniature device. In some embodiments, the world reconstruction component 516 may use a compute budget less than a single Advanced RISC Machine (ARM) core to generate physical world representations in real-time on a non-predefined space such that the remaining compute budget of the single ARM core can be accessed for other uses such as, for example, extracting meshes.

[0158] In some embodiments, the remote data repository 574 may include a digital data storage facility, which may be available through the Internet or other networking configuration in a "cloud" resource configuration. In some embodiments, all data is stored and all computations are performed in the local data processing module 570, allowing fully autonomous use from a remote module. In some embodiments, all data is stored and all or most computations are performed in the remote data repository 574, allowing for a smaller device. A world reconstruction, for example, may be stored in whole or in part in this repository 574.

[0159] In embodiments in which data is stored remotely, and accessible over a network, data may be shared by multiple users of an augmented reality system. For example, user devices may upload their tracking maps to augment a database of environment maps. In some embodiments, the tracking map upload occurs at the end of a user session with a wearable device. In some embodiments, the tracking map uploads may occur continuously, semi-continuously, intermittently, at a pre-defined time, after a pre-defined period from the previous upload, or when triggered by an event. A tracking map uploaded by any user device may be used to expand or improve a previously stored map, whether based on data from that user device or any other user device. Likewise, a persistent map downloaded to a user device may be based on data from that user device or any other user device. In this way, high quality environment maps may be readily available to users to improve their experiences with the AR system.

[0160] In further embodiments, persistent map downloads can be limited and/or avoided based on localization executed on remote resources (e.g., in the cloud). In such configurations, a wearable device or other XR device communicates to the cloud service feature information coupled with pose information (e.g., positioning information for the device at the time the features represented in the feature information were sensed). One or more components of the cloud service may match the feature information to respective stored maps

(e.g., canonical maps) and generates transforms between a tracking map maintained by the XR device and the coordinate system of the canonical map. Each XR device that has its tracking map localized with respect to the canonical map may accurately render virtual content in locations specified with respect to the canonical map based on its own tracking.

[0161] In some embodiments, the local data processing module 570 is operatively coupled to a battery 582. In some embodiments, the battery 582 is a removable power source, such as over the counter batteries. In other embodiments, the battery 582 is a lithium-ion battery. In some embodiments, the battery 582 includes both an internal lithium-ion battery chargeable by the user 560 during non-operation times of the system 580 and removable batteries such that the user 560 may operate the system 580 for longer periods of time without having to be tethered to a power source to charge the lithium-ion battery or having to shut the system 580 off to replace batteries.

[0162] FIG. 5A illustrates a user 530 wearing an AR display system rendering AR content as the user 530 moves through a physical world environment 532 (hereinafter referred to as “environment 532”). The information captured by the AR system along the movement path of the user may be processed into one or more tracking maps. The user 530 positions the AR display system at positions 534, and the AR display system records ambient information of a passable world (e.g., a digital representation of the real objects in the physical world that can be stored and updated with changes to the real objects in the physical world) relative to the positions 534. That information may be stored as poses in combination with images, features, directional audio inputs, or other desired data. The positions 534 are aggregated to data inputs 536, for example, as part of a tracking map, and processed at least by a passable world module 538, which may be implemented, for example, by processing on a remote processing module 572 of FIG. 4. In some embodiments, the passable world module 538 may include the headpose component 514 and the world reconstruction component 516, such that the processed information may indicate the location of objects in the physical world in combination with other information about physical objects used in rendering virtual content.

[0163] The passable world module 538 determines, at least in part, where and how AR content 540 can be placed in the physical world as determined from the data inputs 536. The AR content is “placed” in the physical world by presenting via the user interface both a representation of the physical world and the AR content, with the AR content rendered as if it were interacting with objects in the physical world and the objects in the physical world presented as if the AR content were, when appropriate, obscuring the user’s view of those objects. In some embodiments, the AR content may be placed by appropriately selecting portions of a fixed element 542 (e.g., a table) from a reconstruction (e.g., the reconstruction 518) to determine the shape and position of the AR content 540. As an example, the fixed element may be a table and the virtual content may be positioned such that it appears to be on that table. In some embodiments, the AR content may be placed within structures in a field of view 544, which may be a present field of view or an estimated future field of view. In some embodiments, the AR content may be persisted relative to a model 546 of the physical world (e.g. a mesh).

[0164] As depicted, the fixed element 542 serves as a proxy (e.g. digital copy) for any fixed element within the physical world which may be stored in the passable world module 538 so that the user 530 can perceive content on the fixed element 542 without the system having to map to the fixed element 542 each time the user 530 sees it. The fixed element 542 may, therefore, be a mesh model from a previous modeling session or determined from a separate user but nonetheless stored by the passable world module 538 for future reference by a plurality of users. Therefore, the passable world module 538 may recognize the environment 532 from a previously mapped environment and display AR content without a device of the user 530 mapping all or part of the environment 532 first, saving computation process and cycles and avoiding latency of any rendered AR content.

[0165] The mesh model 546 of the physical world may be created by the AR display system and appropriate surfaces and metrics for interacting and displaying the AR content 540 can be stored by the passable world module 538 for future retrieval by the user 530 or other users without the need to completely or partially recreate the model. In some embodiments, the data inputs 536 are inputs such as geolocation, user identification, and current activity to indicate to the passable world module 538 which fixed element 542 of one or more fixed elements are available, which AR content 540 has last been placed on the fixed element 542, and whether to display that same content (such AR content being “persistent” content regardless of user viewing a particular passable world model).

[0166] Even in embodiments in which objects are considered to be fixed (e.g. a kitchen table), the passable world module 538 may update those objects in a model of the physical world from time to time to account for the possibility of changes in the physical world. The model of fixed objects may be updated with a very low frequency. Other objects in the physical world may be moving or otherwise not regarded as fixed (e.g. kitchen chairs). To render an AR scene with a realistic feel, the AR system may update the position of these non-fixed objects with a much higher frequency than is used to update fixed objects. To enable accurate tracking of all of the objects in the physical world, an AR system may draw information from multiple sensors, including one or more image sensors.

[0167] FIG. 5B is a schematic illustration of a viewing optics assembly 548 and attendant components. In some embodiments, two eye tracking cameras 550, directed toward user eyes 549, detect metrics of the user eyes 549, such as eye shape, eyelid occlusion, pupil direction and glint on the user eyes 549.

[0168] In some embodiments, one of the sensors may be a depth sensor 551, such as a time of flight sensor, emitting signals to the world and detecting reflections of those signals from nearby objects to determine distance to given objects. A depth sensor, for example, may quickly determine whether objects have entered the field of view of the user, either as a result of motion of those objects or a change of pose of the user. However, information about the position of objects in the field of view of the user may alternatively or additionally be collected with other sensors. Depth information, for example, may be obtained from stereoscopic visual image sensors or plenoptic sensors.

[0169] In some embodiments, world cameras 552 record a greater-than-peripheral view to map and/or otherwise create

a model of the environment **532** and detect inputs that may affect AR content. In some embodiments, the world camera **552** and/or camera **553** may be grayscale and/or color image sensors, which may output grayscale and/or color image frames at fixed time intervals. Camera **553** may further capture physical world images within a field of view of the user at a specific time. Pixels of a frame-based image sensor may be sampled repetitively even if their values are unchanged. Each of the world cameras **552**, the camera **553** and the depth sensor **551** have respective fields of view of **554**, **555**, and **556** to collect data from and record a physical world scene, such as the physical world environment **532** depicted in FIG. 34A.

[0170] Inertial measurement units **557** may determine movement and orientation of the viewing optics assembly **548**. In some embodiments, inertial measurement units **557** may provide an output indicating a direction of gravity. In some embodiments, each component is operatively coupled to at least one other component. For example, the depth sensor **551** is operatively coupled to the eye tracking cameras **550** as a confirmation of measured accommodation against actual distance the user eyes **549** are looking at.

[0171] It should be appreciated that a viewing optics assembly **548** may include some of the components illustrated in FIG. 34B and may include components instead of or in addition to the components illustrated. In some embodiments, for example, a viewing optics assembly **548** may include two world camera **552** instead of four. Alternatively or additionally, cameras **552** and **553** need not capture a visible light image of their full field of view. A viewing optics assembly **548** may include other types of components. In some embodiments, a viewing optics assembly **548** may include one or more dynamic vision sensor (DVS), whose pixels may respond asynchronously to relative changes in light intensity exceeding a threshold.

[0172] In some embodiments, a viewing optics assembly **548** may not include the depth sensor **551** based on time of flight information. In some embodiments, for example, a viewing optics assembly **548** may include one or more plenoptic cameras, whose pixels may capture light intensity and an angle of the incoming light, from which depth information can be determined. For example, a plenoptic camera may include an image sensor overlaid with a transmissive diffraction mask (TDM). Alternatively or additionally, a plenoptic camera may include an image sensor containing angle-sensitive pixels and/or phase-detection auto-focus pixels (PDAF) and/or micro-lens array (MLA). Such a sensor may serve as a source of depth information instead of or in addition to depth sensor **551**.

[0173] It also should be appreciated that the configuration of the components in FIG. 5B is provided as an example. A viewing optics assembly **548** may include components with any suitable configuration, which may be set to provide the user with the largest field of view practical for a particular set of components. For example, if a viewing optics assembly **548** has one world camera **552**, the world camera may be placed in a center region of the viewing optics assembly instead of at a side.

[0174] Information from the sensors in viewing optics assembly **548** may be coupled to one or more of processors in the system. The processors may generate data that may be rendered so as to cause the user to perceive virtual content interacting with objects in the physical world. That rendering may be implemented in any suitable way, including

generating image data that depicts both physical and virtual objects. In other embodiments, physical and virtual content may be depicted in one scene by modulating the opacity of a display device that a user looks through at the physical world. The opacity may be controlled so as to create the appearance of the virtual object and also to block the user from seeing objects in the physical world that are occluded by the virtual objects. In some embodiments, the image data may only include virtual content that may be modified such that the virtual content is perceived by a user as realistically interacting with the physical world (e.g. clip content to account for occlusions), when viewed through the user interface.

[0175] The location on the viewing optics assembly **548** at which content is displayed to create the impression of an object at a particular location may depend on the physics of the viewing optics assembly. Additionally, the pose of the user's head with respect to the physical world and the direction in which the user's eyes are looking may impact where in the physical world content displayed at a particular location on the viewing optics assembly content will appear. Sensors as described above may collect this information, and/or supply information from which this information may be calculated, such that a processor receiving sensor inputs may compute where objects should be rendered on the viewing optics assembly **548** to create a desired appearance for the user.

[0176] Regardless of how content is presented to a user, a model of the physical world may be used so that characteristics of the virtual objects, which can be impacted by physical objects, including the shape, position, motion, and visibility of the virtual object, can be correctly computed. In some embodiments, the model may include the reconstruction of a physical world, for example, the reconstruction **518**.

[0177] That model may be created from data collected from sensors on a wearable device of the user. Though, in some embodiments, the model may be created from data collected by multiple users, which may be aggregated in a computing device remote from all of the users (and which may be "in the cloud").

[0178] The model may be created, at least in part, by a world reconstruction system such as, for example, the world reconstruction component **516** of FIG. 3 depicted in more detail in FIG. 6A. The world reconstruction component **516** may include a perception module **660** that may generate, update, and store representations for a portion of the physical world. In some embodiments, the perception module **660** may represent the portion of the physical world within a reconstruction range of the sensors as multiple voxels. Each voxel may correspond to a 3D cube of a predetermined volume in the physical world, and include surface information, indicating whether there is a surface in the volume represented by the voxel. Voxels may be assigned values indicating whether their corresponding volumes have been determined to include surfaces of physical objects, determined to be empty or have not yet been measured with a sensor and so their value is unknown. It should be appreciated that values indicating that voxels that are determined to be empty or unknown need not be explicitly stored, as the values of voxels may be stored in computer memory in any suitable way, including storing no information for voxels that are determined to be empty or unknown.

[0179] In addition to generating information for a persisted world representation, the perception module **660** may identify and output indications of changes in a region around a user of an AR system. Indications of such changes may trigger updates to volumetric data stored as part of the persisted world, or trigger other functions, such as triggering components **604** that generate AR content to update the AR content.

[0180] In some embodiments, the perception module **660** may identify changes based on a signed distance function (SDF) model. The perception module **660** may be configured to receive sensor data such as, for example, depth maps **660a** and headposes **660b**, and then fuse the sensor data into a SDF model **660c**. Depth maps **660a** may provide SDF information directly, and images may be processed to arrive at SDF information. The SDF information represents distance from the sensors used to capture that information. As those sensors may be part of a wearable unit, the SDF information may represent the physical world from the perspective of the wearable unit and therefore the perspective of the user. The headposes **660b** may enable the SDF information to be related to a voxel in the physical world.

[0181] In some embodiments, the perception module **660** may generate, update, and store representations for the portion of the physical world that is within a perception range. The perception range may be determined based, at least in part, on a sensor's reconstruction range, which may be determined based, at least in part, on the limits of a sensor's observation range. As a specific example, an active depth sensor that operates using active IR pulses may operate reliably over a range of distances, creating the observation range of the sensor, which may be from a few centimeters or tens of centimeters to a few meters.

[0182] The world reconstruction component **516** may include additional modules that may interact with the perception module **660**. In some embodiments, a persisted world module **662** may receive representations for the physical world based on data acquired by the perception module **660**. The persisted world module **662** also may include various formats of representations of the physical world. For example, volumetric metadata **662b** such as voxels may be stored as well as meshes **662c** and planes **662d**. In some embodiments, other information, such as depth maps could be saved.

[0183] In some embodiments, representations of the physical world, such as those illustrated in FIG. 6A may provide relatively dense information about the physical world in comparison to sparse maps, such as a tracking map based on feature points as described above.

[0184] In some embodiments, the perception module **660** may include modules that generate representations for the physical world in various formats including, for example, meshes **660d**, planes and semantics **660e**. The representations for the physical world may be stored across local and remote storage mediums. The representations for the physical world may be described in different coordinate frames depending on, for example, the location of the storage medium. For example, a representation for the physical world stored in the device may be described in a coordinate frame local to the device. The representation for the physical world may have a counterpart stored in a cloud. The counterpart in the cloud may be described in a coordinate frame shared by all devices in an XR system.

[0185] In some embodiments, these modules may generate representations based on data within the perception range of one or more sensors at the time the representation is generated as well as data captured at prior times and information in the persisted world module **662**. In some embodiments, these components may operate on depth information captured with a depth sensor. However, the AR system may include vision sensors and may generate such representations by analyzing monocular or binocular vision information.

[0186] In some embodiments, these modules may operate on regions of the physical world. Those modules may be triggered to update a subregion of the physical world, when the perception module **660** detects a change in the physical world in that subregion. Such a change, for example, may be detected by detecting a new surface in the SDF model **660c** or other criteria, such as changing the value of a sufficient number of voxels representing the subregion.

[0187] The world reconstruction component **516** may include components **664** that may receive representations of the physical world from the perception module **660**. Information about the physical world may be pulled by these components according to, for example, a use request from an application. In some embodiments, information may be pushed to the use components, such as via an indication of a change in a pre-identified region or a change of the physical world representation within the perception range. The components **664**, may include, for example, game programs and other components that perform processing for visual occlusion, physics-based interactions, and environment reasoning.

[0188] Responding to the queries from the components **664**, the perception module **660** may send representations for the physical world in one or more formats. For example, when the component **664** indicates that the use is for visual occlusion or physics-based interactions, the perception module **660** may send a representation of surfaces. When the component **664** indicates that the use is for environmental reasoning, the perception module **660** may send meshes, planes and semantics of the physical world.

[0189] In some embodiments, the perception module **660** may include components that format information to provide the component **664**. An example of such a component may be raycasting component **660f**. A use component (e.g., component **664**), for example, may query for information about the physical world from a particular point of view. Raycasting component **660f** may select from one or more representations of the physical world data within a field of view from that point of view.

[0190] As should be appreciated from the foregoing description, the perception module **660**, or another component of an AR system, may process data to create 3D representations of portions of the physical world. Data to be processed may be reduced by culling parts of a 3D reconstruction volume based at last in part on a camera frustum and/or depth image, extracting and persisting plane data, capturing, persisting, and updating 3D reconstruction data in blocks that allow local update while maintaining neighbor consistency, providing occlusion data to applications generating such scenes, where the occlusion data is derived from a combination of one or more depth data sources, and/or performing a multi-stage mesh simplification. The reconstruction may contain data of different levels of sophis-

tication including, for example, raw data such as live depth data, fused volumetric data such as voxels, and computed data such as meshes.

[0191] In some embodiments, components of a passable world model may be distributed, with some portions executing locally on an XR device and some portions executing remotely, such as on a network connected server, or otherwise in the cloud. The allocation of the processing and storage of information between the local XR device and the cloud may impact functionality and user experience of an XR system. For example, reducing processing on a local device by allocating processing to the cloud may enable longer battery life and reduce heat generated on the local device. But, allocating too much processing to the cloud may create undesirable latency that causes an unacceptable user experience.

[0192] FIG. 6B depicts a distributed component architecture **600** configured for spatial computing, according to some embodiments. The distributed component architecture **600** may include a passable world component **602** (e.g., PW **538** in FIG. 5A), a Lumin OS **604**, API's **606**, SDK **608**, and Application **610**. The Lumin OS **604** may include a Linux-based kernel with custom drivers compatible with an XR device. The API's **606** may include application programming interfaces that grant XR applications (e.g., Applications **610**) access to the spatial computing features of an XR device. The SDK **608** may include a software development kit that allows the creation of XR applications.

[0193] One or more components in the architecture **600** may create and maintain a model of a passable world. In this example sensor data is collected on a local device. Processing of that sensor data may be performed in part locally on the XR device and partially in the cloud. PW **538** may include environment maps created based, at least in part, on data captured by AR devices worn by multiple users. During sessions of an AR experience, individual AR devices (such as wearable devices described above in connection with FIG. 4) may create tracking maps, which is one type of map.

[0194] In some embodiments, the device may include components that construct both sparse maps and dense maps. A tracking map may serve as a sparse map and may include headposes of the AR device scanning an environment as well as information about objects detected within that environment at each headpose. Those headposes may be maintained locally for each device. For example, the headpose on each device may be relative to an initial headpose when the device was turned on for its session. As a result, each tracking map may be local to the device creating it and may have its own frame of reference defined by its own local coordinate system. In some embodiments, however, the tracking map on each device may be formed such that one coordinate of its local coordinate system is aligned with the direction of gravity as measured by its sensors, such as inertial measurement unit **557**.

[0195] The dense map may include surface information, which may be represented by a mesh or depth information. Alternatively or additionally, a dense map may include higher level information derived from surface or depth information, such as the location and/or characteristics of planes and/or other objects.

[0196] Creation of the dense maps may be independent of the creation of sparse maps, in some embodiments. The creation of dense maps and sparse maps, for example, may be performed in separate processing pipelines within an AR

system. Separating processing, for example, may enable generation or processing of different types of maps to be performed at different rates. Sparse maps, for example, may be refreshed at a faster rate than dense maps. In some embodiments, however, the processing of dense and sparse maps may be related, even if performed in different pipelines. Changes in the physical world revealed in a sparse map, for example, may trigger updates of a dense map, or vice versa. Further, even if independently created, the maps might be used together. For example, a coordinate system derived from a sparse map may be used to define position and/or orientation of objects in a dense map.

[0197] The sparse map and/or dense map may be persisted for re-use by the same device and/or sharing with other devices. Such persistence may be achieved by storing information in the cloud. The AR device may send the tracking map to a cloud to, for example, merge with environment maps selected from persisted maps previously stored in the cloud. In some embodiments, the selected persisted maps may be sent from the cloud to the AR device for merging. In some embodiments, the persisted maps may be oriented with respect to one or more persistent coordinate frames. Such maps may serve as canonical maps, as they can be used by any of multiple devices. In some embodiments, a model of a passable world may comprise or be created from one or more canonical maps. Devices, even though they perform some operations based on a coordinate frame local to the device, may nonetheless use the canonical map by determining a transformation between their coordinate frame local to the device and the canonical map.

[0198] A canonical map may originate as a tracking map (TM) (e.g., TM **1102** in FIG. 31A), which may be promoted to a canonical map. The canonical map may be persisted such that devices that access the canonical map may, once determining a transformation between their local coordinate system and a coordinate system of the canonical map, use the information in the canonical map to determine locations of objects represented in the canonical map in the physical world around the device. In some embodiments, a TM may be a headpose sparse map created by an XR device. In some embodiments, the canonical map may be created when an XR device sends one or more TMs to a cloud server for merging with additional TMs captured by the XR device at a different time or by other XR devices.

[0199] In embodiments in which tracking maps are formed on local devices with one coordinate of a local coordinate frame aligned with gravity, this orientation with respect to gravity may be preserved upon creation of a canonical map. For example, when a tracking map that is submitted for merging does not overlap with any previously stored map, that tracking map may be promoted to a canonical map. Other tracking maps, which may also have an orientation relative to gravity, may be subsequently merged with that canonical map. The merging may be done so as to ensure that the resulting canonical map retains its orientation relative to gravity. Two maps, for example, may not be merged, regardless of correspondence of feature points in those maps, if coordinates of each map aligned with gravity do not align with each other with a sufficiently close tolerance.

[0200] The canonical maps, or other maps, may provide information about the portions of the physical world represented by the data processed to create respective maps. FIG. 7 depicts an exemplary tracking map **700**, according to some embodiments. The tracking map **700** may provide a floor

plan 706 of physical objects in a corresponding physical world, represented by points 702. In some embodiments, a map point 702 may represent a feature of a physical object that may include multiple features. For example, each corner of a table may be a feature that is represented by a point on a map. The features may be derived from processing images, such as may be acquired with the sensors of a wearable device in an augmented reality system. The features, for example, may be derived by processing an image frame output by a sensor to identify features based on large gradients in the image or other suitable criteria. Further processing may limit the number of features in each frame. For example, processing may select features that likely represent persistent objects. One or more heuristics may be applied for this selection.

[0201] The tracking map 700 may include data on points 702 collected by a device. For each image frame with data points included in a tracking map, a pose may be stored. The pose may represent the orientation from which the image frame was captured, such that the feature points within each image frame may be spatially correlated. The pose may be determined by positioning information, such as may be derived from the sensors, such as an IMU sensor, on the wearable device. Alternatively or additionally, the pose may be determined from matching image frames to other image frames that depict overlapping portions of the physical world. By finding such positional correlation, which may be accomplished by matching subsets of features points in two frames, the relative pose between the two frames may be computed. A relative pose may be adequate for a tracking map, as the map may be relative to a coordinate system local to a device established based on the initial pose of the device when construction of the tracking map was initiated.

[0202] Not all of the feature points and image frames collected by a device may be retained as part of the tracking map, as much of the information collected with the sensors is likely to be redundant. Rather, only certain frames may be added to the map. Those frames may be selected based on one or more criteria, such as degree of overlap with image frames already in the map, the number of new features they contain or a quality metric for the features in the frame. Image frames not added to the tracking map may be discarded or may be used to revise the location of features. As a further alternative, all or most of the image frames, represented as a set of features may be retained, but a subset of those frames may be designated as key frames, which are used for further processing.

[0203] The key frames may be processed to produce keyrigs 704. The key frames may be processed to produce three dimensional sets of feature points and saved as keyrigs 704. Such processing may entail, for example, comparing image frames derived simultaneously from two cameras to stereoscopically determine the 3D position of feature points. Metadata may be associated with these keyframes and/or keyrigs, such as poses.

[0204] The environment maps may have any of multiple formats depending on, for example, the storage locations of an environment map including, for example, local storage of AR devices and remote storage. For example, a map in remote storage may have higher resolution than a map in local storage on a wearable device where memory is limited. To send a higher resolution map from remote storage to local storage, the map may be down sampled or otherwise converted to an appropriate format, such as by reducing the

number of poses per area of the physical world stored in the map and/or the number of feature points stored for each pose. In some embodiments, a slice or portion of a high resolution map from remote storage may be sent to local storage, where the slice or portion is not down sampled.

[0205] A database of environment maps may be updated as new tracking maps are created. To determine which of a potentially very large number of environment maps in a database is to be updated, updating may include efficiently selecting one or more environment maps stored in the database relevant to the new tracking map. The selected one or more environment maps may be ranked by relevance and one or more of the highest ranking maps may be selected for processing to merge higher ranked selected environment maps with the new tracking map to create one or more updated environment maps. When a new tracking map represents a portion of the physical world for which there is no preexisting environment map to update, that tracking map may be stored in the database as a new environment map.

View Independent Display

[0206] Described herein are methods and apparatus for providing virtual contents using an XR system, independent of locations of eyes viewing the virtual content. Conventionally, a virtual content is re-rendered upon any motion of the displaying system. For example, if a user wearing a display system views a virtual representation of a three-dimensional (3D) object on the display and walks around the area where the 3D object appears, the 3D object should be re-rendered for each viewpoint such that the user has the perception that he or she is walking around an object that occupies real space. However, the re-rendering consumes significant computational resources of a system and causes artifacts due to latency.

[0207] The inventors have recognized and appreciated that headpose (e.g., the location and orientation of a user wearing an XR system) may be used to render a virtual content independent of eye rotations within a head of the user. In some embodiments, dynamic maps of a scene may be generated based on multiple coordinate frames in real space across one or more sessions such that virtual contents interacting with the dynamic maps may be rendered robustly, independent of eye rotations within the head of the user and/or independent of sensor deformations caused by, for example, heat generated during high-speed, computation-intensive operation. In some embodiments, the configuration of multiple coordinate frames may enable a first XR device worn by a first user and a second XR device worn by a second user to recognize a common location in a scene. In some embodiments, the configuration of multiple coordinate frames may enable users wearing XR devices to view a virtual content in a same location of a scene.

[0208] In some embodiments, a tracking map may be built in a world coordinate frame, which may have a world origin. The world origin may be the first pose of an XR device when the XR device is powered on. The world origin may be aligned to gravity such that a developer of an XR application can get gravity alignment without extra work. Different tracking maps may be built in different world coordinate frames because the tracking maps may be captured by a same XR device at different sessions and/or different XR devices worn by different users. In some embodiments, a session of an XR device may span from powering on to

powering off the device. In some embodiments, an XR device may have a head coordinate frame, which may have a head origin. The head origin may be the current pose of an XR device when an image is taken. The difference between headpose of a world coordinate frame and of a head coordinate frame may be used to estimate a tracking route.

[0209] In some embodiments, an XR device may have a camera coordinate frame, which may have a camera origin. The camera origin may be the current pose of one or more sensors of an XR device. The inventors have recognized and appreciated that the configuration of a camera coordinate frame enables robust displaying virtual contents independent of eye rotation within a head of a user. This configuration also enables robust displaying of virtual contents independent of sensor deformation due to, for example, heat generated during operation.

[0210] In some embodiments, an XR device may have a head unit with a head-mountable frame that a user can secure to their head and may include two waveguides, one in front of each eye of the user. The waveguides may be transparent so that ambient light from real-world objects can transmit through the waveguides and the user can see the real-world objects. Each waveguide may transmit projected light from a projector to a respective eye of the user. The projected light may form an image on the retina of the eye. The retina of the eye thus receives the ambient light and the projected light. The user may simultaneously see real-world objects and one or more virtual objects that are created by the projected light. In some embodiments, XR devices may have sensors that detect real-world objects around a user. These sensors may, for example, be cameras that capture images that may be processed to identify the locations of real-world objects.

[0211] In some embodiments, an XR system may assign a coordinate frame to a virtual content, as opposed to attaching the virtual content in a world coordinate frame. Such configuration enables a virtual content to be described without regard to where it is rendered for a user, but it may be attached to a more persistent frame position such as a persistent coordinate frame (PCF) described in relation to, for example, FIGS. 14-20C, to be rendered in a specified location. When the locations of the objects change, the XR device may detect the changes in the environment map and determine movement of the head unit worn by the user relative to real-world objects.

[0212] FIG. 8 illustrates a user experiencing virtual content, as rendered by an XR system 10, in a physical environment, according to some embodiments. The XR system may include a first XR device 12.1 that is worn by a first user 14.1, a network 18 and a server 20. The user 14.1 is in a physical environment with a real object in the form of a table 16.

[0213] In the illustrated example, the first XR device 12.1 includes a head unit 22, a belt pack 24 and a cable connection 26. The first user 14.1 secures the head unit 22 to their head and the belt pack 24 remotely from the head unit 22 on their waist. The cable connection 26 connects the head unit 22 to the belt pack 24. The head unit 22 includes technologies that are used to display a virtual object or objects to the first user 14.1 while the first user 14.1 is permitted to see real objects such as the table 16. The belt pack 24 includes primarily processing and communications capabilities of the first XR device 12.1. In some embodiments, the processing and communication capabilities may reside entirely or par-

tially in the head unit 22 such that the belt pack 24 may be removed or may be located in another device such as a backpack.

[0214] In the illustrated example, the belt pack 24 is connected via a wireless connection to the network 18. The server 20 is connected to the network 18 and holds data representative of local content. The belt pack 24 downloads the data representing the local content from the server 20 via the network 18. The belt pack 24 provides the data via the cable connection 26 to the head unit 22. The head unit 22 may include a display that has a light source, for example, a laser light source or a light emitting diode (LED), and a waveguide that guides the light.

[0215] In some embodiments, the first user 14.1 may mount the head unit 22 to their head and the belt pack 24 to their waist. The belt pack 24 may download image data representing virtual content over the network 18 from the server 20. The first user 14.1 may see the table 16 through a display of the head unit 22. A projector forming part of the head unit 22 may receive the image data from the belt pack 24 and generate light based on the image data. The light may travel through one or more of the waveguides forming part of the display of the head unit 22. The light may then leave the waveguide and propagates onto a retina of an eye of the first user 14.1. The projector may generate the light in a pattern that is replicated on a retina of the eye of the first user 14.1. The light that falls on the retina of the eye of the first user 14.1 may have a selected field of depth so that the first user 14.1 perceives an image at a preselected depth behind the waveguide. In addition, both eyes of the first user 14.1 may receive slightly different images so that a brain of the first user 14.1 perceives a three-dimensional image or images at selected distances from the head unit 22. In the illustrated example, the first user 14.1 perceives a virtual content 28 above the table 16. The proportions of the virtual content 28 and its location and distance from the first user 14.1 are determined by the data representing the virtual content 28 and various coordinate frames that are used to display the virtual content 28 to the first user 14.1.

[0216] In the illustrated example, the virtual content 28 is not visible from the perspective of the drawing and is visible to the first user 14.1 through using the first XR device 12.1. The virtual content 28 may initially reside as data structures within vision data and algorithms in the belt pack 24. The data structures may then manifest themselves as light when the projectors of the head unit 22 generate light based on the data structures. It should be appreciated that although the virtual content 28 has no existence in three-dimensional space in front of the first user 14.1, the virtual content 28 is still represented in FIG. 1 in three-dimensional space for illustration of what a wearer of head unit 22 perceives. The visualization of computer data in three-dimensional space may be used in this description to illustrate how the data structures that facilitate the renderings are perceived by one or more users relate to one another within the data structures in the belt pack 24.

[0217] FIG. 9 illustrates components of the first XR device 12.1, according to some embodiments. The first XR device 12.1 may include the head unit 22, and various components forming part of the vision data and algorithms including, for example, a rendering engine 30, various coordinate systems 32, various origin and destination coordinate frames 34, and various origin to destination coordinate frame transformers 36. The various coordinate systems may be based on intrin-

sics of to the XR device or may be determined by reference to other information, such as a persistent pose or a persistent coordinate system, as described herein.

[0218] The head unit 22 may include a head-mountable frame 40, a display system 42, a real object detection camera 44, a movement tracking camera 46, and an inertial measurement unit 48.

[0219] The head-mountable frame 40 may have a shape that is securable to the head of the first user 14.1 in FIG. 8. The display system 42, real object detection camera 44, movement tracking camera 46, and inertial measurement unit 48 may be mounted to the head-mountable frame 40 and therefore move together with the head-mountable frame 40.

[0220] The coordinate systems 32 may include a local data system 52, a world frame system 54, a head frame system 56, and a camera frame system 58.

[0221] The local data system 52 may include a data channel 62, a local frame determining routine 64 and a local frame storing instruction 66. The data channel 62 may be an internal software routine, a hardware component such as an external cable or a radio frequency receiver, or a hybrid component such as a port that is opened up. The data channel 62 may be configured to receive image data 68 representing a virtual content.

[0222] The local frame determining routine 64 may be connected to the data channel 62. The local frame determining routine 64 may be configured to determine a local coordinate frame 70. In some embodiments, the local frame determining routine may determine the local coordinate frame based on real world objects or real world locations. In some embodiments, the local coordinate frame may be based on a top edge relative to a bottom edge of a browser window, head or feet of a character, a node on an outer surface of a prism or bounding box that encloses the virtual content, or any other suitable location to place a coordinate frame that defines a facing direction of a virtual content and a location (e.g. a node, such as a placement node or PCF node) with which to place the virtual content, etc.

[0223] The local frame storing instruction 66 may be connected to the local frame determining routine 64. One skilled in the art will understand that software modules and routines are “connected” to one another through subroutines, calls, etc. The local frame storing instruction 66 may store the local coordinate frame 70 as a local coordinate frame 72 within the origin and destination coordinate frames 34. In some embodiments, the origin and destination coordinate frames 34 may be one or more coordinate frames that may be manipulated or transformed in order for a virtual content to persist between sessions. In some embodiments, a session may be the period of time between a boot-up and shut-down of an XR device. Two sessions may be two start-up and shut-down periods for a single XR device, or may be a start-up and shut-down for two different XR devices.

[0224] In some embodiments, the origin and destination coordinate frames 34 may be the coordinate frames involved in one or more transformations required in order for a first user's XR device and a second user's XR device to recognize a common location. In some embodiments, the destination coordinate frame may be the output of a series of computations and transformations applied to the target coordinate frame in order for a first and second user to view a virtual content in the same location.

[0225] The rendering engine 30 may be connected to the data channel 62. The rendering engine 30 may receive the

image data 68 from the data channel 62 such that the rendering engine 30 may render virtual content based, at least in part, on the image data 68.

[0226] The display system 42 may be connected to the rendering engine 30. The display system 42 may include components that transform the image data 68 into visible light. The visible light may form two patterns, one for each eye. The visible light may enter eyes of the first user 14.1 in FIG. 8 and may be detected on retinas of the eyes of the first user 14.1.

[0227] The real object detection camera 44 may include one or more cameras that may capture images from different sides of the head-mountable frame 40. The movement tracking camera 46 may include one or more cameras that capture images on sides of the head-mountable frame 40. One set of one or more cameras may be used instead of the two sets of one or more cameras representing the real object detection camera(s) 44 and the movement tracking camera(s) 46. In some embodiments, the cameras 44, 46 may capture images. As described above these cameras may collect data that is used to construct a tracking map.

[0228] The inertial measurement unit 48 may include a number of devices that are used to detect movement of the head unit 22. The inertial measurement unit 48 may include a gravitation sensor, one or more accelerometers and one or more gyroscopes. The sensors of the inertial measurement unit 48, in combination, track movement of the head unit 22 in at least three orthogonal directions and about at least three orthogonal axes.

[0229] In the illustrated example, the world frame system 54 includes a world surface determining routine 78, a world frame determining routine 80, and a world frame storing instruction 82. The world surface determining routine 78 is connected to the real object detection camera 44. The world surface determining routine 78 receives images and/or key frames based on the images that are captured by the real object detection camera 44 and processes the images to identify surfaces in the images. A depth sensor (not shown) may determine distances to the surfaces. The surfaces are thus represented by data in three dimensions including their sizes, shapes, and distances from the real object detection camera.

[0230] In some embodiments, a world coordinate frame 84 may be based on the origin at the initialization of the headpose session. In some embodiments, the world coordinate frame may be located where the device was booted up, or could be somewhere new if headpose was lost during the boot session. In some embodiments, the world coordinate frame may be the origin at the start of a headpose session.

[0231] In the illustrated example, the world frame determining routine 80 is connected to the world surface determining routine 78 and determines a world coordinate frame 84 based on the locations of the surfaces as determined by the world surface determining routine 78. The world frame storing instruction 82 is connected to the world frame determining routine 80 to receive the world coordinate frame 84 from the world frame determining routine 80. The world frame storing instruction 82 stores the world coordinate frame 84 as a world coordinate frame 86 within the origin and destination coordinate frames 34.

[0232] The head frame system 56 may include a head frame determining routine 90 and a head frame storing instruction 92. The head frame determining routine 90 may be connected to the movement tracking camera 46 and the

inertial measurement unit **48**. The head frame determining routine **90** may use data from the movement tracking camera **46** and the inertial measurement unit **48** to calculate a head coordinate frame **94**. For example, the inertial measurement unit **48** may have a gravitation sensor that determines the direction of gravitational force relative to the head unit **22**. The movement tracking camera **46** may continually capture images that are used by the head frame determining routine **90** to refine the head coordinate frame **94**. The head unit **22** moves when the first user **14.1** in FIG. 8 moves their head. The movement tracking camera **46** and the inertial measurement unit **48** may continuously provide data to the head frame determining routine **90** so that the head frame determining routine **90** can update the head coordinate frame **94**. [0233] The head frame storing instruction **92** may be connected to the head frame determining routine **90** to receive the head coordinate frame **94** from the head frame determining routine **90**. The head frame storing instruction **92** may store the head coordinate frame **94** as a head coordinate frame **96** among the origin and destination coordinate frames **34**. The head frame storing instruction **92** may repeatedly store the updated head coordinate frame **94** as the head coordinate frame **96** when the head frame determining routine **90** recalculates the head coordinate frame **94**. In some embodiments, the head coordinate frame may be the location of the wearable XR device **12.1** relative to the local coordinate frame **72**.

[0234] The camera frame system **58** may include camera intrinsics **98**. The camera intrinsics **98** may include dimensions of the head unit **22** that are features of its design and manufacture. The camera intrinsics **98** may be used to calculate a camera coordinate frame **100** that is stored within the origin and destination coordinate frames **34**.

[0235] In some embodiments, the camera coordinate frame **100** may include all pupil positions of a left eye of the first user **14.1** in FIG. 8. When the left eye moves from left to right or up and down, the pupil positions of the left eye are located within the camera coordinate frame **100**. In addition, the pupil positions of a right eye are located within a camera coordinate frame **100** for the right eye. In some embodiments, the camera coordinate frame **100** may include the location of the camera relative to the local coordinate frame when an image is taken.

[0236] The origin to destination coordinate frame transformers **36** may include a local-to-world coordinate transformer **104**, a world-to-head coordinate transformer **106**, and a head-to-camera coordinate transformer **108**. The local-to-world coordinate transformer **104** may receive the local coordinate frame **72** and transform the local coordinate frame **72** to the world coordinate frame **86**. The transformation of the local coordinate frame **72** to the world coordinate frame **86** may be represented as a local coordinate frame transformed to world coordinate frame **110** within the world coordinate frame **86**.

[0237] The world-to-head coordinate transformer **106** may transform from the world coordinate frame **86** to the head coordinate frame **96**. The world-to-head coordinate transformer **106** may transform the local coordinate frame transformed to world coordinate frame **110** to the head coordinate frame **96**. The transformation may be represented as a local coordinate frame transformed to head coordinate frame **112** within the head coordinate frame **96**.

[0238] The head-to-camera coordinate transformer **108** may transform from the head coordinate frame **96** to the

camera coordinate frame **100**. The head-to-camera coordinate transformer **108** may transform the local coordinate frame transformed to head coordinate frame **112** to a local coordinate frame transformed to camera coordinate frame **114** within the camera coordinate frame **100**. The local coordinate frame transformed to camera coordinate frame **114** may be entered into the rendering engine **30**. The rendering engine **30** may render the image data **68** representing the local content **28** based on the local coordinate frame transformed to camera coordinate frame **114**.

[0239] FIG. 10 is a spatial representation of the various origin and destination coordinate frames **34**. The local coordinate frame **72**, world coordinate frame **86**, head coordinate frame **96**, and camera coordinate frame **100** are represented in the figure. In some embodiments, the local coordinate frame associated with the XR content **28** may have a position and rotation (e.g. may provide a node and facing direction) relative to a local and/or world coordinate frame and/or PCF when the virtual content is placed in the real world so the virtual content may be viewed by the user. Each camera may have its own camera coordinate frame **100** encompassing all pupil positions of one eye. Reference numerals **104A** and **106A** represent the transformations that are made by the local-to-world coordinate transformer **104**, world-to-head coordinate transformer **106**, and head-to-camera coordinate transformer **108** in FIG. 9, respectively.

[0240] FIG. 11 depicts a camera render protocol for transforming from a head coordinate frame to a camera coordinate frame, according to some embodiments. In the illustrated example, a pupil for a single eye moves from position A to B. A virtual object that is meant to appear stationary will project onto a depth plane at one of the two positions A or B depending on the position of the pupil (assuming that the camera is configured to use a pupil-based coordinate frame). As a result, using a pupil coordinate frame transformed to a head coordinate frame will cause jitter in a stationary virtual object as the eye moves from position A to position B. This situation is referred to as view dependent display or projection.

[0241] As depicted in FIG. 12, a camera coordinate frame (e.g., CR) is positioned and encompasses all pupil positions and object projection will now be consistent regardless of pupil positions A and B. The head coordinate frame transforms to the CR frame, which is referred to as view independent display or projection. An image reprojection may be applied to the virtual content to account for a change in eye position, however, as the rendering is still in the same position, jitter is minimized.

[0242] FIG. 13 illustrates the display system **42** in more detail. The display system **42** includes a stereoscopic analyzer **144** that is connected to the rendering engine **30** and forms part of the vision data and algorithms.

[0243] The display system **42** further includes left and right projectors **166A** and **166B** and left and right waveguides **170A** and **170B**. The left and right projectors **166A** and **166B** are connected to power supplies. Each projector **166A** and **166B** has a respective input for image data to be provided to the respective projector **166A** or **166B**. The respective projector **166A** or **166B**, when powered, generates light in two-dimensional patterns and emanates the light therefrom. The left and right waveguides **170A** and **170B** are positioned to receive light from the left and right projectors **166A** and **166B**, respectively. The left and right waveguides **170A** and **170B** are transparent waveguides.

[0244] In use, a user mounts the head mountable frame 40 to their head. Components of the head mountable frame 40 may, for example, include a strap (not shown) that wraps around the back of the head of the user. The left and right waveguides 170A and 170B are then located in front of left and right eyes 220A and 220B of the user.

[0245] The rendering engine 30 enters the image data that it receives into the stereoscopic analyzer 144. The image data is three-dimensional image data of the local content 28 in FIG. 8. The image data is projected onto a plurality of virtual planes. The stereoscopic analyzer 144 analyzes the image data to determine left and right image data sets based on the image data for projection onto each depth plane. The left and right image data sets are data sets that represent two-dimensional images that are projected in three-dimensions to give the user a perception of a depth.

[0246] The stereoscopic analyzer 144 enters the left and right image data sets into the left and right projectors 166A and 166B. The left and right projectors 166A and 166B then create left and right light patterns. The components of the display system 42 are shown in plan view, although it should be understood that the left and right patterns are two-dimensional patterns when shown in front elevation view. Each light pattern includes a plurality of pixels. For purposes of illustration, light rays 224A and 226A from two of the pixels are shown leaving the left projector 166A and entering the left waveguide 170A. The light rays 224A and 226A reflect from sides of the left waveguide 170A. It is shown that the light rays 224A and 226A propagate through internal reflection from left to right within the left waveguide 170A, although it should be understood that the light rays 224A and 226A also propagate in a direction into the paper using refractory and reflective systems.

[0247] The light rays 224A and 226A exit the left light waveguide 170A through a pupil 228A and then enter a left eye 220A through a pupil 230A of the left eye 220A. The light rays 224A and 226A then fall on a retina 232A of the left eye 220A. In this manner, the left light pattern falls on the retina 232A of the left eye 220A. The user is given the perception that the pixels that are formed on the retina 232A are pixels 234A and 236A that the user perceives to be at some distance on a side of the left waveguide 170A opposing the left eye 220A. Depth perception is created by manipulating the focal length of the light.

[0248] In a similar manner, the stereoscopic analyzer 144 enters the right image data set into the right projector 166B. The right projector 166B transmits the right light pattern, which is represented by pixels in the form of light rays 224B and 226B. The light rays 224B and 226B reflect within the right waveguide 170B and exit through a pupil 228B. The light rays 224B and 226B then enter through a pupil 230B of the right eye 220B and fall on a retina 232B of a right eye 220B. The pixels of the light rays 224B and 226B are perceived as pixels 134B and 236B behind the right waveguide 170B.

[0249] The patterns that are created on the retinas 232A and 232B are individually perceived as left and right images. The left and right images differ slightly from one another due to the functioning of the stereoscopic analyzer 144. The left and right images are perceived in a mind of the user as a three-dimensional rendering.

[0250] As mentioned, the left and right waveguides 170A and 170B are transparent. Light from a real-life object such as the table 16 on a side of the left and right waveguides

170A and 170B opposing the eyes 220A and 220B can project through the left and right waveguides 170A and 170B and fall on the retinas 232A and 232B.

Persistent Coordinate Frame (PCF)

[0251] Described herein are methods and apparatus for providing spatial persistence across user instances within a shared space. Without spatial persistence, virtual content placed in the physical world by a user in a session may not exist or may be misplaced in the user's view in a different session. Without spatial persistence, virtual content placed in the physical world by one user may not exist or may be out of place in a second user's view, even if the second user is intended to be sharing an experience of the same physical space with the first user.

[0252] The inventors have recognized and appreciated that spatial persistence may be provided through persistent coordinate frames (PCFs). A PCF may be defined based on one or more points, representing features recognized in the physical world (e.g., corners, edges). The features may be selected such that they are likely to be the same from a user instance to another user instance of an XR system.

[0253] Further, drift during tracking, which causes the computed tracking path (e.g., camera trajectory) to deviate from the actual tracking path, can cause the location of virtual content, when rendered with respect to a local map that is based solely on a tracking map to appear out of place. A tracking map for the space may be refined to correct the drifts as an XR device collects more information of the scene overtime. However, if virtual content is placed on a real object before a map refinement and saved with respect to the world coordinate frame of the device derived from the tracking map, the virtual content may appear displaced, as if the real object has been moved during the map refinement. PCFs may be updated according to map refinement because the PCFs are defined based on the features and are updated as the features move during map refinements.

[0254] A PCF may comprise six degrees of freedom with translations and rotations relative to a map coordinate system. A PCF may be stored in a local and/or remote storage medium. The translations and rotations of a PCF may be computed relative to a map coordinate system depending on, for example, the storage location. For example, a PCF used locally by a device may have translations and rotations relative to a world coordinate frame of the device. A PCF in the cloud may have translations and rotations relative to a canonical coordinate frame of a canonical map.

[0255] PCFs may provide a sparse representation of the physical world, providing less than all of the available information about the physical world, such that they may be efficiently processed and transferred. Techniques for processing persistent spatial information may include creating dynamic maps based on one or more coordinate systems in real space across one or more sessions, generating persistent coordinate frames (PCF) over the sparse maps, which may be exposed to XR applications via, for example, an application programming interface (API).

[0256] FIG. 14 is a block diagram illustrating the creation of a persistent coordinate frame (PCF) and the attachment of XR content to the PCF, according to some embodiments. Each block may represent digital information stored in a computer memory. In the case of applications 1180, the data may represent computer-executable instructions. In the case of virtual content 1170, the digital information may define a

virtual object, as specified by the application 1180, for example. In the case of the other boxes, the digital information may characterize some aspect of the physical world.

[0257] In the illustrated embodiment, one or more PCFs are created from images captured with sensors on a wearable device. In the embodiment of FIG. 14, the sensors are visual image cameras. These cameras may be the same cameras used for forming a tracking map. Accordingly, some of the processing suggested by FIG. 14 may be performed as part of updating a tracking map. However, FIG. 14 illustrates that information that provides persistence is generated in addition to the tracking map.

[0258] In order to derive a 3D PCF, two images 1110 from two cameras mounted to a wearable device in a configuration that enables stereoscopic image analysis are processed together. FIG. 14 illustrates an Image 1 and an Image 2, each derived from one of the cameras. A single image from each camera is illustrated for simplicity. However, each camera may output a stream of image frames and the processing illustrated in FIG. 14 may be performed for multiple image frames in the stream.

[0259] Accordingly, Image 1 and Image 2 may each be one frame in a sequence of image frames. Processing as depicted in FIG. 14 may be repeated on successive image frames in the sequence until image frames containing feature points providing a suitable image from which to form persistent spatial information is processed. Alternatively or additionally, the processing of FIG. 14 might be repeated as a user moves such that the user is no longer close enough to a previously identified PCF to reliably use that PCF for determining positions with respect to the physical world. For example, an XR system may maintain a current PCF for a user. When that distance exceeds a threshold, the system may switch to a new current PCF, closer to the user, which may be generated according to the process of FIG. 14, using image frames acquired in the user's current location.

[0260] Even when generating a single PCF, a stream of image frames may be processed to identify image frames depicting content in the physical world that is likely stable and can be readily identified by a device in the vicinity of the region of the physical world depicted in the image frame. In the embodiment of FIG. 14, this processing begins with the identification of features 1120 in the image. Features may be identified, for example, by finding locations of gradients in the image above a threshold or other characteristics, which may correspond to a corner of an object, for example. In the embodiment illustrated, the features are points, but other recognizable features, such as edges, may alternatively or additionally be used.

[0261] In the embodiment illustrated, a fixed number, N, of features 1120 are selected for further processing. Those feature points may be selected based on one or more criteria, such as magnitude of the gradient, or proximity to other feature points. Alternatively or additionally, the feature points may be selected heuristically, such as based on characteristics that suggest the feature points are persistent. For example, heuristics may be defined based on the characteristics of feature points that likely correspond to a corner of a window or a door or a large piece of furniture. Such heuristics may take into account the feature point itself and what surrounds it. As a specific example, the number of feature points per image may be between 100 and 500 or between 150 and 250, such as 200.

[0262] Regardless of the number of feature points selected, descriptors 1130 may be computed for the feature points. In this example, a descriptor is computed for each selected feature point, but a descriptor may be computed for groups of feature points or for a subset of the feature points or for all features within an image. The descriptor characterizes a feature point such that feature points representing the same object in the physical world are assigned similar descriptors. The descriptors may facilitate alignment of two frames, such as may occur when one map is localized with respect to another. Rather than searching for a relative orientation of the frames that minimizes the distance between feature points of the two images, an initial alignment of the two frames may be made by identifying feature points with similar descriptors. Alignment of the image frames may be based on aligning points with similar descriptors, which may entail less processing than computing an alignment of all the feature points in the images.

[0263] The descriptors may be computed as a mapping of the feature points or, in some embodiments a mapping of a patch of an image around a feature point, to a descriptor. The descriptor may be a numeric quantity. U.S. patent application Ser. No. 16/190,948 describes computing descriptors for feature points and is hereby incorporated herein by reference in its entirety.

[0264] In the example of FIG. 14, a descriptor 1130 is computed for each feature point in each image frame. Based on the descriptors and/or the feature points and/or the image itself, the image frame may be identified as a key frame 1140. In the embodiment illustrated, a key frame is an image frame meeting certain criteria that is then selected for further processing. In making a tracking map, for example, image frames that add meaningful information to the map may be selected as key frames that are integrated into the map. On the other hand, image frames that substantially overlap a region for which an image frame has already been integrated into the map may be discarded such that they do not become key frames. Alternatively or additionally, key frames may be selected based on the number and/or type of feature points in the image frame. In the embodiment of FIG. 14, key frames 1150 selected for inclusion in a tracking map may also be treated as key frames for determining a PCF, but different or additional criteria for selecting key frames for generation of a PCF may be used.

[0265] Though FIG. 14 shows that a key frame is used for further processing, information acquired from an image may be processed in other forms. For example, the feature points, such as in a key rig, may alternatively or additionally be processed. Moreover, though a key frame is described as being derived from a single image frame, it is not necessary that there be a one to one relationship between a key frame and an acquired image frame. A key frame, for example, may be acquired from multiple image frames, such as by stitching together or aggregating the image frames such that only features appearing in multiple images are retained in the key frame.

[0266] A key frame may include image information and/or metadata associated with the image information. In some embodiments, images captured by the cameras 44, 46 (FIG. 9) may be computed into one or more key frames (e.g., key frames 1, 2). In some embodiments, a key frame may include a camera pose. In some embodiments, a key frame may include one or more camera images captured at the camera pose. In some embodiments, an XR system may

determine a portion of the camera images captured at the camera pose as not useful and thus not include the portion in a key frame. Therefore, using key frames to align new images with earlier knowledge of a scene reduces the use of computational resource of the XR system. In some embodiments, a key frame may include an image, and/or image data, at a location with a direction/angle. In some embodiments, a key frame may include a location and a direction from which one or more map points may be observed. In some embodiments, a key frame may include a coordinate frame with an ID. U.S. patent application Ser. No. 15/877,359 describes key frames and is hereby incorporated herein by reference in its entirety.

[0267] Some or all of the key frames **1140** may be selected for further processing, such as the generation of a persistent pose **1150** for the key frame. The selection may be based on the characteristics of all, or a subset of, the feature points in the image frame. Those characteristics may be determined from processing the descriptors, features and/or image frame, itself. As a specific example, the selection may be based on a cluster of feature points identified as likely to relate to a persistent object.

[0268] Each key frame is associated with a pose of the camera at which that key frame was acquired. For key frames selected for processing into a persistent pose, that pose information may be saved along with other metadata about the key frame, such as a WiFi fingerprint and/or GPS coordinates at the time of acquisition and/or at the location of acquisition.

[0269] The persistent poses are a source of information that a device may use to orient itself relative to previously acquired information about the physical world. For example, if the key frame from which a persistent pose was created is incorporated into a map of the physical world, a device may orient itself relative to that persistent pose using a sufficient number of feature points in the key frame that are associated with the persistent pose. The device may align a current image that it takes of its surroundings to the persistent pose. This alignment may be based on matching the current image to the image **1110**, the features **1120**, and/or the descriptors **1130** that gave rise to the persistent pose, or any subset of that image or those features or descriptors. In some embodiments, the current image frame that is matched to the persistent pose may be another key frame that has been incorporated into the device's tracking map.

[0270] Information about a persistent pose may be stored in a format that facilitates sharing among multiple applications, which may be executing on the same or different devices. In the example of FIG. 14, some or all of the persistent poses may be reflected as a persistent coordinate frames (PCF) **1160**. Like a persistent pose, a PCF may be associated with a map and may comprise a set of features, or other information, that a device can use to determine its orientation with respect to that PCF. The PCF may include a transformation that defines its transformation with respect to the origin of its map, such that, by correlating its position to a PCF, the device can determine its position with respect to any objects in the physical world reflected in the map.

[0271] As the PCF provides a mechanism for determining locations with respect to the physical objects, an application, such as applications **1180**, may define positions of virtual objects with respect to one or more PCFs, which serve as anchors for the virtual content **1170**. FIG. 14 illustrates, for example, that App **1** has associated its virtual content **2** with

PCF **1.2**. Likewise, App **2** has associated its virtual content **3** with PCF **1.2**. App **1** is also shown associating its virtual content **1** to PCF **4.5**, and App **2** is shown associating its virtual content **4** with PCF **3**. In some embodiments, PCF **3** may be based on Image **3** (not shown), and PCF **4.5** may be based on Image **4** and Image **5** (not shown) analogously to how PCF **1.2** is based on Image **1** and Image **2**. When rendering this virtual content, a device may apply one or more transformations to compute information, such as the location of the virtual content with respect to the display of the device and/or the location of physical objects with respect to the desired location of the virtual content. Using the PCFs as reference may simplify such computations.

[0272] In some embodiments, a persistent pose may be a coordinate location and/or direction that has one or more associated key frames. In some embodiments, a persistent pose may be automatically created after the user has traveled a certain distance, e.g., three meters. In some embodiments, the persistent poses may act as reference points during localization. In some embodiments, the persistent poses may be stored in a passable world (e.g., the passable world module **538**).

[0273] In some embodiments, a new PCF may be determined based on a pre-defined distance allowed between adjacent PCFs. In some embodiments, one or more persistent poses may be computed into a PCF when a user travels a pre-determined distance, e.g. five meters. In some embodiments, PCFs may be associated with one or more world coordinate frames and/or canonical coordinate frames, e.g., in the passable world. In some embodiments, PCFs may be stored in a local and/or remote database depending on, for example, security settings.

[0274] FIG. 15 illustrates a method **4700** of establishing and using a persistence coordinate frame, according to some embodiments. The method **4700** may start from capturing (Act **4702**) images (e.g., Image **1** and Image **2** in FIG. 14) about a scene using one or more sensors of an XR device. Multiple cameras may be used and one camera may generate multiple images, for example, in a stream.

[0275] The method **4700** may include extracting (4704) interest points (e.g., map points **702** in FIG. 7, features **1120** in FIG. 14) from the captured images, generating (Act **4706**) descriptors (e.g., descriptors **1130** in FIG. 14) for the extracted interest points, and generating (Act **4708**) key frames (e.g., key frames **1140**) based on the descriptors. In some embodiments, the method may compare interest points in the key frames, and form pairs of key frames that share a predetermined amount of interest points. The method may reconstruct parts of the physical world using individual pairs of key frames. Mapped parts of the physical world may be saved as 3D features (e.g., keyrig **704** in FIG. 7). In some embodiments, a selected portion of the pairs of key frames may be used to build 3D features. In some embodiments, results of the mapping may be selectively saved. Key frames not used for building 3D features may be associated with the 3D features through poses, for example, representing distances between key frames with a covariance matrix between poses of keyframes. In some embodiments, pairs of key frames may be selected to build the 3D features such that distances between each two of the build 3D features are within a predetermined distance, which may be determined to balance the amount of computation needed and the level of accuracy of a resulting model. Such approaches enable providing a model of the physical world with the amount of

data that is suitable for efficient and accurate computation with an XR system. In some embodiments, a covariance matrix of two images may include covariances between poses of the two images (e.g., six degrees of freedom).

[0276] The method 4700 may include generating (Act 4710) persistent poses based on the key frames. In some embodiments, the method may include generating the persistent poses based on the 3D features reconstructed from pairs of key frames. In some embodiments, a persistent pose may be attached to a 3D feature. In some embodiments, the persistent pose may include a pose of a key frame used to construct the 3D feature. In some embodiments, the persistent pose may include an average pose of key frames used to construct the 3D feature. In some embodiments, persistent poses may be generated such that distances between neighboring persistent poses are within a predetermined value, for example, in the range of one meter to five meters, any value in between, or any other suitable value. In some embodiments, the distances between neighboring persistent poses may be represented by a covariance matrix of the neighboring persistent poses.

[0277] The method 4700 may include generating (Act 4712) PCFs based on the persistent poses. In some embodiments, a PCF may be attached to a 3D feature. In some embodiments, a PCF may be associated with one or more persistent poses. In some embodiments, a PCF may include a pose of one of the associated persistent poses. In some embodiments, a PCF may include an average pose of the poses of the associated persistent poses. In some embodiments, PCFs may be generated such that distances between neighboring PCFs are within a predetermined value, for example, in the range of three meters to ten meters, any value in between, or any other suitable value. In some embodiments, the distances between neighboring PCFs may be represented by a covariance matrix of the neighboring PCFs. In some embodiments, PCFs may be exposed to XR applications via, for example, an application programming interface (API) such that the XR applications can access a model of the physical world through the PCFs without accessing the model itself.

[0278] The method 4700 may include associating (Act 4714) image data of a virtual object to be displayed by the XR device to at least one of the PCFs. In some embodiments, the method may include computing translations and orientations of the virtual object with respect to the associated PCF. It should be appreciated that it is not necessary to associate a virtual object to a PCF generated by the device placing the virtual object. For example, a device may retrieve saved PCFs in a canonical map in a cloud and associate a virtual object to a retrieved PCF. It should be appreciated that the virtual object may move with the associated PCF as the PCF is adjusted overtime.

[0279] FIG. 16 illustrates the first XR device 12.1 and vision data and algorithms of a second XR device 12.2 and the server 20, according to some embodiments. The components illustrated in FIG. 16 may operate to perform some or all of the operations associated with generating, updating, and/or using spatial information, such as persistent poses, persistent coordinate frames, tracking maps, or canonical maps, as described herein. Although not illustrated, the first XR device 12.1 may be configured the same as the second XR device 12.2. The server 20 may have a map storing routine 118, a canonical map 120, a map transmitter 122, and a map merge algorithm 124.

[0280] The second XR device 12.2, which may be in the same scene as the first XR device 12.1, may include a persistent coordinate frame (PCF) integration unit 1300, an application 1302 that generates the image data 68 that may be used to render a virtual object, and a frame embedding generator 308 (See FIG. 21). In some embodiments, a map download system 126, PCF identification system 128, Map 2, localization module 130, canonical map incorporator 132, canonical map 133, and map publisher 136 may be grouped into a passable world unit 1304. The PCF integration unit 1300 may be connected to the passable world unit 1304 and other components of the second XR device 12.2 to allow for the retrieval, generation, use, upload, and download of PCFs.

[0281] A map, comprising PCFs, may enable more persistence in a changing world. In some embodiments, localizing a tracking map including, for example, matching features for images, may include selecting features that represent persistent content from the map constituted by PCFs, which enables fast matching and/or localizing. For example, a world where people move into and out of the scene and objects such as doors move relative to the scene, requires less storage space and transmission rates, and enables the use of individual PCFs and their relationships relative to one another (e.g., integrated constellation of PCFs) to map a scene.

[0282] In some embodiments, the PCF integration unit 1300 may include PCFs 1306 that were previously stored in a data store on a storage unit of the second XR device 12.2, a PCF tracker 1308, a persistent pose acquirer 1310, a PCF checker 1312, a PCF generation system 1314, a coordinate frame calculator 1316, a persistent pose calculator 1318, and three transformers, including a tracking map and persistent pose transformer 1320, a persistent pose and PCF transformer 1322, and a PCF and image data transformer 1324.

[0283] In some embodiments, the PCF tracker 1308 may have an on-prompt and an off-prompt that are selectable by the application 1302. The application 1302 may be executable by a processor of the second XR device 12.2 to, for example, display a virtual content. The application 1302 may have a call that switches the PCF tracker 1308 on via the on-prompt. The PCF tracker 1308 may generate PCFs when the PCF tracker 1308 is switched on. The application 1302 may have a subsequent call that can switch the PCF tracker 1308 off via the off-prompt. The PCF tracker 1308 terminates PCF generation when the PCF tracker 1308 is switched off.

[0284] In some embodiments, the server 20 may include a plurality of persistent poses 1332 and a plurality of PCFs 1330 that have previously been saved in association with a canonical map 120. The map transmitter 122 may transmit the canonical map 120 together with the persistent poses 1332 and/or the PCFs 1330 to the second XR device 12.2. The persistent poses 1332 and PCFs 1330 may be stored in association with the canonical map 133 on the second XR device 12.2. When Map 2 localizes to the canonical map 133, the persistent poses 1332 and the PCFs 1330 may be stored in association with Map 2.

[0285] In some embodiments, the persistent pose acquirer 1310 may acquire the persistent poses for Map 2. The PCF checker 1312 may be connected to the persistent pose acquirer 1310. The PCF checker 1312 may retrieve PCFs from the PCFs 1306 based on the persistent poses retrieved by the persistent pose acquirer 1310. The PCFs retrieved by

the PCF checker **1312** may form an initial group of PCFs that are used for image display based on PCFs.

[0286] In some embodiments, the application **1302** may require additional PCFs to be generated. For example, if a user moves to an area that has not previously been mapped, the application **1302** may switch the PCF tracker **1308** on. The PCF generation system **1314** may be connected to the PCF tracker **1308** and begin to generate PCFs based on Map **2** as Map **2** begins to expand. The PCFs generated by the PCF generation system **1314** may form a second group of PCFs that may be used for PCF-based image display.

[0287] The coordinate frame calculator **1316** may be connected to the PCF checker **1312**. After the PCF checker **1312** retrieved PCFs, the coordinate frame calculator **1316** may invoke the head coordinate frame **96** to determine a head-pose of the second XR device **12.2**. The coordinate frame calculator **1316** may also invoke the persistent pose calculator **1318**. The persistent pose calculator **1318** may be directly or indirectly connected to the frame embedding generator **308**. In some embodiments, an image/frame may be designated a key frame after a threshold distance from the previous key frame, e.g. 3 meters, is traveled. The persistent pose calculator **1318** may generate a persistent pose based on a plurality, for example three, key frames. In some embodiments, the persistent pose may be essentially an average of the coordinate frames of the plurality of key frames.

[0288] The tracking map and persistent pose transformer **1320** may be connected to Map **2** and the persistent pose calculator **1318**. The tracking map and persistent pose transformer **1320** may transform Map **2** to the persistent pose to determine the persistent pose at an origin relative to Map **2**.

[0289] The persistent pose and PCF transformer **1322** may be connected to the tracking map and persistent pose transformer **1320** and further to the PCF checker **1312** and the PCF generation system **1314**. The persistent pose and PCF transformer **1322** may transform the persistent pose (to which the tracking map has been transformed) to the PCFs from the PCF checker **1312** and the PCF generation system **1314** to determine the PCFs relative to the persistent pose.

[0290] The PCF and image data transformer **1324** may be connected to the persistent pose and PCF transformer **1322** and to the data channel **62**. The PCF and image data transformer **1324** transforms the PCFs to the image data **68**. The rendering engine **30** may be connected to the PCF and image data transformer **1324** to display the image data **68** to the user relative to the PCFs.

[0291] The PCF integration unit **1300** may store the additional PCFs that are generated with the PCF generation system **1314** within the PCFs **1306**. The PCFs **1306** may be stored relative to persistent poses. The map publisher **136** may retrieve the PCFs **1306** and the persistent poses associated with the PCFs **1306** when the map publisher **136** transmits Map **2** to the server **20**, the map publisher **136** also transmits the PCFs and persistent poses associated with Map **2** to the server **20**. When the map storing routine **118** of the server **20** stores Map **2**, the map storing routine **118** may also store the persistent poses and PCFs generated by the second viewing device **12.2**. The map merge algorithm **124** may create the canonical map **120** with the persistent poses and PCFs of Map **2** associated with the canonical map **120** and stored within the persistent poses **1332** and PCFs **1330**, respectively.

[0292] The first XR device **12.1** may include a PCF integration unit similar to the PCF integration unit **1300** of the second XR device **12.2**. When the map transmitter **122** transmits the canonical map **120** to the first XR device **12.1**, the map transmitter **122** may transmit the persistent poses **1332** and PCFs **1330** associated with the canonical map **120** and originating from the second XR device **12.2**. The first XR device **12.1** may store the PCFs and the persistent poses within a data store on a storage device of the first XR device **12.1**. The first XR device **12.1** may then make use of the persistent poses and the PCFs originating from the second XR device **12.2** for image display relative to the PCFs. Additionally or alternatively, the first XR device **12.1** may retrieve, generate, make use, upload, and download PCFs and persistent poses in a manner similar to the second XR device **12.2** as described above.

[0293] In the illustrated example, the first XR device **12.1** generates a local tracking map (referred to hereinafter as “Map **1**”) and the map storing routine **118** receives Map **1** from the first XR device **12.1**. The map storing routine **118** then stores Map **1** on a storage device of the server **20** as the canonical map **120**.

[0294] The second XR device **12.2** includes a map download system **126**, an anchor identification system **128**, a localization module **130**, a canonical map incorporator **132**, a local content position system **134**, and a map publisher **136**.

[0295] In use, the map transmitter **122** sends the canonical map **120** to the second XR device **12.2** and the map download system **126** downloads and stores the canonical map **120** as a canonical map **133** from the server **20**.

[0296] The anchor identification system **128** is connected to the world surface determining routine **78**. The anchor identification system **128** identifies anchors based on objects detected by the world surface determining routine **78**. The anchor identification system **128** generates a second map (Map **2**) using the anchors. As indicated by the cycle **138**, the anchor identification system **128** continues to identify anchors and continues to update Map **2**. The locations of the anchors are recorded as three-dimensional data based on data provided by the world surface determining routine **78**. The world surface determining routine **78** receives images from the real object detection camera **44** and depth data from depth sensors **135** to determine the locations of surfaces and their relative distance from the depth sensors **135**.

[0297] The localization module **130** is connected to the canonical map **133** and Map **2**. The localization module **130** repeatedly attempts to localize Map **2** to the canonical map **133**. The canonical map incorporator **132** is connected to the canonical map **133** and Map **2**. When the localization module **130** localizes Map **2** to the canonical map **133**, the canonical map incorporator **132** incorporates the canonical map **133** into anchors of Map **2**. Map **2** is then updated with missing data that is included in the canonical map.

[0298] The local content position system **134** is connected to Map **2**. The local content position system **134** may, for example, be a system wherein a user can locate local content in a particular location within a world coordinate frame. The local content then attaches itself to one anchor of Map **2**. The local-to-world coordinate transformer **104** transforms the local coordinate frame to the world coordinate frame based on the settings of the local content position system **134**. The

functioning of the rendering engine **30**, display system **42**, and data channel **62** have been described with reference to FIG. 2.

[0299] The map publisher **136** uploads Map **2** to the server **20**. The map storing routine **118** of the server **20** then stores Map **2** within a storage medium of the server **20**.

[0300] The map merge algorithm **124** merges Map **2** with the canonical map **120**. When more than two maps, for example, three or four maps relating to the same or adjacent regions of the physical world, have been stored, the map merge algorithm **124** merges all the maps into the canonical map **120** to render a new canonical map **120**. The map transmitter **122** then transmits the new canonical map **120** to any and all devices **12.1** and **12.2** that are in an area represented by the new canonical map **120**. When the devices **12.1** and **12.2** localize their respective maps to the canonical map **120**, the canonical map **120** becomes the promoted map.

[0301] FIG. 17 illustrates an example of generating key frames for a map of a scene, according to some embodiments. In the illustrated example, a first key frame, KF₁, is generated for a door on a left wall of the room. A second key frame, KF₂, is generated for an area in a corner where a floor, the left wall, and a right wall of the room meet. A third key frame, KF₃, is generated for an area of a window on the right wall of the room. A fourth key frame, KF₄, is generated for an area at a far end of a rug on a floor of the wall. A fifth key frame, KF₅, is generated for an area of the rug closest to the user.

[0302] FIG. 18 illustrates an example of generating persistent poses for the map of FIG. 17, according to some embodiments. In some embodiments, a new persistent pose is created when the device measures a threshold distance traveled, and/or when an application requests a new persistent pose (PP). In some embodiments, the threshold distance may be 3 meters, 5 meters, 20 meters, or any other suitable distance. Selecting a smaller threshold distance (e.g., 1 m) may result in an increase in compute load since a larger number of PPs may be created and managed compared to larger threshold distances. Selecting a larger threshold distance (e.g. 40 m) may result in increased virtual content placement error since a smaller number of PPs would be created, which would result in fewer PCFs being created, which means the virtual content attached to the PCF could be a relatively large distance (e.g. 30 m) away from the PCF, and error increases with increasing distance from a PCF to the virtual content.

[0303] In some embodiments, a PP may be created at the start of a new session. This initial PP may be thought of as zero, and can be visualized as the center of a circle that has a radius equal to the threshold distance. When the device reaches the perimeter of the circle, and, in some embodiments, an application requests a new PP, a new PP may be placed at the current location of the device (at the threshold distance). In some embodiments, a new PP will not be created at the threshold distance if the device is able to find an existing PP within the threshold distance from the device's new position. In some embodiments, when a new PP (e.g., PP**1150** in FIG. 14) is created, the device attaches one or more of the closest key frames to the PP. In some embodiments, the location of the PP relative to the key frames may be based on the location of the device at the time

a PP is created. In some embodiments, a PP will not be created when the device travels a threshold distance unless an application requests a PP.

[0304] In some embodiments, an application may request a PCF from the device when the application has virtual content to display to the user. The PCF request from the application may trigger a PP request, and a new PP would be created after the device travels the threshold distance. FIG. 18 illustrates a first persistent pose PP₁ which may have the closest key frames, (e.g. KF₁, KF₂, and KF₃) attached by, for example, computing relative poses between the key frames to the persistent pose. FIG. 18 also illustrates a second persistent pose PP₂ which may have the closest key frames (e.g. KF₄ and KF₅) attached.

[0305] FIG. 19 illustrates an example of generating a PCF for the map of FIG. 17, according to some embodiments. In the illustrated example, a PCF 1 may include PP₁ and PP₂. As described above, the PCF may be used for displaying image data relative to the PCF. In some embodiments, each PCF may have coordinates in another coordinate frame (e.g., a world coordinate frame) and a PCF descriptor, for example, uniquely identifying the PCF. In some embodiments, the PCF descriptor may be computed based on feature descriptors of features in frames associated with the PCF. In some embodiments, various constellations of PCFs may be combined to represent the real world in a persistent manner and that requires less data and less transmission of data.

[0306] FIGS. 20A to 20C are schematic diagrams illustrating an example of establishing and using a persistent coordinate frame. FIG. 20A shows two users **4802A**, **4802B** with respective local tracking maps **4804A**, **4804B** that have not localized to a canonical map. The origins **4806A**, **4806B** for individual users are depicted by the coordinate system (e.g., a world coordinate system) in their respective areas. These origins of each tracking map may be local to each user, as the origins are dependent on the orientation of their respective devices when tracking was initiated.

[0307] As the sensors of the user device scan the environment, the device may capture images that, as described above in connection with FIG. 14, may contain features representing persistent objects such that those images may be classified as key frames, from which a persistent pose may be created. In this example, the tracking map **4802A** includes a persistent pose (PP) **4808A**; the tracking **4802B** includes a PP **4808B**.

[0308] Also as described above in connection with FIG. 14, some of the PP's may be classified as PCFs which are used to determine the orientation of virtual content for rendering it to the user. FIG. 20B shows that XR devices worn by respective users **4802A**, **4802B** may create local PCFs **4810A**, **4810B** based on the PP **4808A**, **4808B**. FIG. 20C shows that persistent content **4812A**, **4812B** (e.g., a virtual content) may be attached to the PCFs **4810A**, **4810B** by respective XR devices.

[0309] In this example, virtual content may have a virtual content coordinate frame, that may be used by an application generating virtual content, regardless of how the virtual content should be displayed. The virtual content, for example, may be specified as surfaces, such as triangles of a mesh, at particular locations and angles with respect to the virtual content coordinate frame. To render that virtual

content to a user, the locations of those surfaces may be determined with respect to the user that is to perceive the virtual content.

[0310] Attaching virtual content to the PCFs may simplify the computation involved in determining locations of the virtual content with respect to the user. The location of the virtual content with respect to a user may be determined by applying a series of transformations. Some of those transformations may change, and may be updated frequently. Others of those transformations may be stable and may be updated in frequently or not at all. Regardless, the transformations may be applied with relatively low computational burden such that the location of the virtual content can be updated with respect to the user frequently, providing a realistic appearance to the rendered virtual content.

[0311] In the example of FIGS. 20A-20C, user 1's device has a coordinate system that can be related to the coordinate system that defines the origin of the map by the transformation rig1_T_w1 . User 2's device has a similar transformation rig2_T_w2 . These transformations may be expressed as six degrees of transformation, specifying translation and rotation to align the device coordinate systems with the map coordinate systems. In some embodiments, the transformation may be expressed as two separate transformations, one specifying translation and the other specifying rotation. Accordingly, it should be appreciated that the transformations may be expressed in a form that simplifies computation or otherwise provides an advantage.

[0312] Transformations between the origins of the tracking maps and the PCFs identified by the respective user devices are expressed as pcf1_T_w1 and pcf2_T_w2 . In this example the PCF and the PP are identical, such that the same transformation also characterizes the PP's.

[0313] The location of the user device with respect to the PCF can therefore be computed by the serial application of these transformations, such as $\text{rig1_T_pcf1} = (\text{rig1_T_w1}) * (\text{pcf1_T_w1})$.

[0314] As shown in FIG. 20C, the virtual content is located with respect to the PCFs, with a transformation of obj1_T_pcf1 . This transformation may be set by an application generating the virtual content that may receive information from a world reconstruction system describing physical objects with respect to the PCF. To render the virtual content to the user, a transformation to the coordinate system of the user's device is computed, which may be computed by relating the virtual content coordinate frame to the origin of the tracking map through the transformation $\text{obj1_t_w1} = (\text{obj1_T_pcf1}) * (\text{pcf1_T_w1})$. That transformation may then be related to the user's device through further transformation rig1_T_w1 .

[0315] The location of the virtual content may change, based on output from an application generating the virtual content. When that changes, the end-to-end transformation, from a source coordinate system to a destination coordinate system, may be recomputed. Additionally, the location and/or headpose of the user may change as the user moves. As a result, the transformation rig1_T_w1 may change, as would any end-to-end transformation that depends on the location or headpose of the user.

[0316] The transformation rig1_T_w1 may be updated with motion of the user based on tracking the position of the user with respect to stationary objects in the physical world. Such tracking may be performed by a headphone tracking component processing a sequence of images, as described

above, or other component of the system. Such updates may be made by determining pose of the user with respect to a stationary frame of reference, such as a PP.

[0317] In some embodiments, the location and orientation of a user device may be determined relative to the nearest persistent pose, or, in this example, a PCF, as the PP is used as a PCF. Such a determination may be made by identifying in current images captured with sensors on the device feature points that characterize the PP. Using image processing techniques, such as stereoscopic image analysis, the location of the device with respect to those feature points may be determined. From this data, the system could calculate the change in transformation associated with the user's motions based on the relationship $\text{rig1_T_pcf1} = (\text{rig1_T_w1}) * (\text{pcf1_T_w1})$.

[0318] A system may determine and apply transformations in an order that is computationally efficient. For example, the need to compute rig1_T_w1 from a measurement yielding rig1_T_pcf1 might be avoided by tracking both user pose and defining the location of virtual content relative to the PP or a PCF built on a persistent pose. In this way the transformation from a source coordinate system of the virtual content to the destination coordinate system of the user's device may be based on the measured transformation according to the expression $(\text{rig1_T_pcf1}) * (\text{obj1_t_pcf1})$, with the first transformation being measured by the system and the latter transformation being supplied by an application specifying virtual content for rendering. In embodiments in which the virtual content is positioned with respect to the origin of the map, the end-to-end transformation may relate the virtual object coordinate system to the PCF coordinate system based on a further transformation between the map coordinates and the PCF coordinates. In embodiments in which the virtual content is positioned with respect to a different PP or PCF than the one against which user position is being tracked, a transformation between the two may be applied. Such a transformation may be fixed and may be determined, for example, from a map in which both appear.

[0319] A transform-based approach may be implemented, for example, in a device with components that process sensor data to build a tracking map. As part of that process, those components may identify feature points that may be used as persistent poses, which in turn may be turned into PCFs. Those components may limit the number of persistent poses generated for the map, to provide a suitable spacing between persistent poses, while allowing the user, regardless of location in the physical environment, to be close enough to a persistent pose location to accurately compute the user's pose, as described above in connection with FIGS. 17-19. As the closest persistent pose to a user is updated, as a result of user movement, refinements to the tracking map or other causes, any of the transformations that are used to compute the location of virtual content relative to the user that depend on the location of the PP (or PCF if being used) may be updated and stored for use, at least until the user moves away from that persistent pose. Nonetheless, by computing and storing transformations, the computational burden each time the location of virtual content is update may be relatively low, such that it may be performed with relatively low latency.

[0320] FIGS. 20A-20C illustrate positioning with respect to a tracking map, and each device had its own tracking map. However, transformations may be generated with respect to

any map coordinate system. Persistence of content across user sessions of an XR system may be achieved by using a persistent map. Shared experiences of users may also be facilitated by using a map to which multiple user devices may be oriented.

[0321] In some embodiments, described in greater detail below, the location of virtual content may be specified in relation to coordinates in a canonical map, formatted such that any of multiple devices may use the map. Each device might maintain a tracking map and may determine the change of pose of the user with respect to the tracking map. In this example, a transformation between the tracking map and the canonical map may be determined through a process of “localization”—which may be performed by matching structures in the tracking map (such as one or more persistent poses) to one or more structures of the canonical map (such as one or more PCFs).

[0322] Described in greater detail below are techniques for creating and using canonical maps in this way.

Deep Key Frame

[0323] Techniques as described herein rely on comparison of image frames. For example, to establish the position of a device with respect to a tracking map, a new image may be captured with sensors worn by the user and an XR system may search, in a set of images that were used to create the tracking map, images that share at least a predetermined amount of interest points with the new image. As an example of another scenario involving comparisons of image frames, a tracking map might be localized to a canonical map by first finding image frames associated with a persistent pose in the tracking map that is similar to an image frame associated with a PCF in the canonical map. Alternatively, a transformation between two canonical maps may be computed by first finding similar image frames in the two maps.

[0324] Deep key frames provide a way to reduce the amount of processing required to identify similar image frames. For example, in some embodiments, the comparison may be between image features in a new 2D image (e.g., “2D features”) and 3D features in the map. Such a comparison may be made in any suitable way, such as by projecting the 3D images into a 2D plane. A conventional method such as Bag of Words (BoW) searches the 2D features of a new image in a database including all 2D features in a map, which may require significant computing resources especially when a map represents a large area. The conventional method then locates the images that share at least one of the 2D features with the new image, which may include images that are not useful for locating meaningful 3D features in the map. The conventional method then locates 3D features that are not meaningful with respect to the 2D features in the new image.

[0325] The inventors have recognized and appreciated techniques to retrieve images in the map using less memory resource (e.g., a quarter of the memory resource used by BoW), higher efficiency (e.g., 2.5 ms processing time for each key frame, 100 μ s for comparing against 500 key frames), and higher accuracy (e.g., 20% better retrieval recall than BoW for 1024 dimensional model, 5% better retrieval recall than BoW for 256 dimensional model).

[0326] To reduce computation, a descriptor may be computed for an image frame that may be used to compare an image frame to other image frames. The descriptors may be stored instead of or in addition to the image frames and

feature points. In a map in which persistent poses and/or PCFs may be generated from image frames, the descriptor of the image frame or frames from which each persistent pose or PCF was generated may be stored as part of the persistent pose and/or PCF.

[0327] In some embodiments, the descriptor may be computed as a function of feature points in the image frame. In some embodiments, a neural network is configured to compute a unique frame descriptor to represent an image. The image may have a resolution higher than 1 Megabyte such that enough details of a 3D environment within a field-of-view of a device worn by a user is captured in the image. The frame descriptor may be much shorter, such as a string of numbers, for example, in the range of 128 Bytes to 512 Bytes or any number in between.

[0328] In some embodiments, the neural network is trained such that the computed frame descriptors indicate similarity between images. Images in a map may be located by identifying, in a database comprising images used to generate the map, the nearest images that may have frame descriptors within a predetermined distance to a frame descriptor for a new image. In some embodiments, the distances between images may be represented by a difference between the frame descriptors of the two images.

[0329] FIG. 21 is a block diagram illustrating a system for generating a descriptor for an individual image, according to some embodiments. In the illustrated example, a frame embedding generator 308 is shown. The frame embedding generator 308, in some embodiments, may be used within the server 20, but may alternatively or additionally execute in whole or in part within one of the XR devices 12.1 and 12.2, or any other device processing images for comparison to other images.

[0330] In some embodiments, the frame embedding generator may be configured to generate a reduced data representation of an image from an initial size (e.g., 76,800 bytes) to a final size (e.g., 256 bytes) that is nonetheless indicative of the content in the image despite a reduced size. In some embodiments, the frame embedding generator may be used to generate a data representation for an image which may be a key frame or a frame used in other ways. In some embodiments, the frame embedding generator 308 may be configured to convert an image at a particular location and orientation into a unique string of numbers (e.g., 256 bytes). In the illustrated example, an image 320 taken by an XR device may be processed by feature extractor 324 to detect interest points 322 in the image 320. Interest points may be or may not be derived from feature points identified as described above for features 1120 (FIG. 14) or as otherwise described herein. In some embodiments, interest points may be represented by descriptors as described above for descriptors 1130 (FIG. 14), which may be generated using a deep sparse feature method. In some embodiments, each interest point 322 may be represented by a string of numbers (e.g., 32 bytes). There may, for example, be n features (e.g., 100) and each feature is represented by a string of 32 bytes.

[0331] In some embodiments, the frame embedding generator 308 may include a neural network 326. The neural network 326 may include a multi-layer perceptron unit 312 and a maximum (max) pool unit 314. In some embodiments, the multi-layer perceptron (MLP) unit 312 may comprise a multi-layer perceptron, which may be trained. In some embodiments, the interest points 322 (e.g., descriptors for the interest points) may be reduced by the multi-layer

perceptron **312**, and may output as weighted combinations **310** of the descriptors. For example, the MLP may reduce n features to m feature that is less than n features.

[0332] In some embodiments, the MLP unit **312** may be configured to perform matrix multiplication. The multi-layer perceptron unit **312** receives the plurality of interest points **322** of an image **320** and converts each interest point to a respective string of numbers (e.g., 256). For example, there may be 100 features and each feature may be represented by a string of 256 numbers. A matrix, in this example, may be created having 100 horizontal rows and 256 vertical columns. Each row may have a series of 256 numbers that vary in magnitude with some being smaller and others being larger. In some embodiments, the output of the MLP may be an $n \times 256$ matrix, where n represents the number of interest points extracted from the image. In some embodiments, the output of the MLP may be an $m \times 256$ matrix, where m is the number of interest points reduced from n.

[0333] In some embodiments, the MLP **312** may have a training phase, during which model parameters for the MLP are determined, and a use phase. In some embodiments, the MLP may be trained as illustrated in FIG. 25. The input training data may comprise data in sets of three, the set of three comprising 1) a query image, 2) a positive sample, and 3) a negative sample. The query image may be considered the reference image.

[0334] In some embodiments, the positive sample may comprise an image that is similar to the query image. For example, in some embodiments, similar may be having the same object in both the query and positive sample image but viewed from a different angle. In some embodiments, similar may be having the same object in both the query and positive sample images but having the object shifted (e.g., left, right, up, down) relative to the other image.

[0335] In some embodiments, the negative sample may comprise an image that is dissimilar to the query image. For example, in some embodiments, a dissimilar image may not contain any objects that are prominent in the query image or may contain only a small portion of a prominent object in the query image (e.g., <10%, 1%). A similar image, in contrast, may have most of an object (e.g. >50%, or >75%) in the query image, for example.

[0336] In some embodiments, interest points may be extracted from the images in the input training data and may be converted to feature descriptors. These descriptors may be computed both for the training images as shown in FIG. 25 and for extracted features in operation of frame embedding generator **308** of FIG. 21. In some embodiments, a deep sparse feature (DSF) process may be used to generate the descriptors (e.g., DSF descriptors) as described in U.S. patent application Ser. No. 16/190,948. In some embodiments, DSF descriptors are $n \times 32$ dimension. The descriptors may then be passed through the model/MLP to create a 256 byte output. In some embodiments, the model/MLP may have the same structure as MLP **312** such that once the model parameters are set through training, the resulting trained MLP may be used as MLP **312**.

[0337] In some embodiments, the feature descriptors (e.g., the 256 byte output from the MLP model) may then be sent to a triplet margin loss module (which may only be used during the training phase, not during use phase of the MLP neural network). In some embodiments, the triplet margin loss module may be configured to select parameters for the model so as to reduce the difference between the 256 byte

output from the query image and the 256 byte output from the positive sample, and to increase the difference between the 256 byte output from the query image and the 256 byte output from the negative sample. In some embodiments, the training phase may comprise feeding a plurality of triplet input images into the learning process to determine model parameters. This training process may continue, for example, until the differences for positive images is minimized and the difference for negative images is maximized or until other suitable exit criteria are reached.

[0338] Referring back to FIG. 21, the frame embedding generator **308** may include a pooling layer, here illustrated as maximum (max) pool unit **314**. The max pool unit **314** may analyze each column to determine a maximum number in the respective column. The max pool unit **314** may combine the maximum value of each column of numbers of the output matrix of the MLP **312** into a global feature string **316** of, for example, 256 numbers. It should be appreciated that images processed in XR systems might, desirably, have high-resolution frames, with potentially millions of pixels. The global feature string **316** is a relatively small number that takes up relatively little memory and is easily searchable compared to an image (e.g., with a resolution higher than 1 Megabyte). It is thus possible to search for images without analyzing each original frame from the camera and it is also cheaper to store 256 bytes instead of complete frames.

[0339] FIG. 22 is a flow chart illustrating a method **2200** of computing an image descriptor, according to some embodiments. The method **2200** may start from receiving (Act **2202**) a plurality of images captured by an XR device worn by a user. In some embodiments, the method **2200** may include determining (Act **2204**) one or more key frames from the plurality of images. In some embodiments, Act **2204** may be skipped and/or may occur after step **2210** instead.

[0340] The method **2200** may include identifying (Act **2206**) one or more interest points in the plurality of images with an artificial neural network, and computing (Act **2208**) feature descriptors for individual interest points with the artificial neural network. The method may include computing (Act **2210**), for each image, a frame descriptor to represent the image based, at least in part, on the computed feature descriptors for the identified interest points in the image with the artificial neural network.

[0341] FIG. 23 is a flow chart illustrating a method **2300** of localization using image descriptors, according to some embodiments. In this example, a new image frame, depicting the current location of the XR device may be compared to image frames stored in connection with points in a map (such as a persistent pose or a PCF as described above). The method **2300** may start from receiving (Act **2302**) a new image captured by an XR device worn by a user. The method **2300** may include identifying (Act **2304**) one or more nearest key frames in a database comprising key frames used to generate one or more maps. In some embodiments, a nearest key frame may be identified based on coarse spatial information and/or previously determined spatial information. For example, coarse spatial information may indicate that the XR device is in a geographic region represented by a 50 m \times 50 m area of a map. Image matching may be performed only for points within that area. As another example, based on tracking, the XR system may know that an XR device was previously proximate a first persistent pose in the map and was moving in a direction of a second

persistent pose in the map. That second persistent pose may be considered the nearest persistent pose and the key frame stored with it may be regarded as the nearest key frame. Alternatively or additionally, other metadata, such as GPS data or WiFi fingerprints, may be used to select a nearest key frame or set of nearest key frames.

[0342] Regardless of how the nearest key frames are selected, frame descriptors may be used to determine whether the new image matches any of the frames selected as being associated with a nearby persistent pose. The determination may be made by comparing a frame descriptor of the new image with frame descriptors of the closest key frames, or a subset of key frames in the database selected in any other suitable way, and selecting key frames with frame descriptors that are within a predetermined distance of the frame descriptor of the new image. In some embodiments, a distance between two frame descriptors may be computed by obtaining the difference between two strings of numbers that may represent the two frame descriptors. In embodiments in which the strings are processed as strings of multiple quantities, the difference may be computed as a vector difference.

[0343] Once a matching image frame is identified, the orientation of the XR device relative to that image frame may be determined. The method 2300 may include performing (Act 2306) feature matching against 3D features in the maps that correspond to the identified nearest key frames, and computing (Act 2308) pose of the device worn by the user based on the feature matching results. In this way, the computationally intensive matching of features points in two images may be performed for as few as one image that has already been determined to be a likely match for the new image.

[0344] FIG. 24 is a flow chart illustrating a method 2400 of training a neural network, according to some embodiments. The method 2400 may start from generating (Act 2402) a dataset comprising a plurality of image sets. Each of the plurality of image sets may include a query image, a positive sample image, and a negative sample image. In some embodiments, the plurality of image sets may include synthetic recording pairs configured to, for example, teach the neural network basic information such as shapes. In some embodiments, the plurality of image sets may include real recording pairs, which may be recorded from a physical world.

[0345] In some embodiments, inliers may be computed by fitting a fundamental matrix between two images. In some embodiments, sparse overlap may be computed as the intersection over union (IoU) of interest points seen in both images. In some embodiments, a positive sample may include at least twenty interest points, serving as inliers, that are the same as in the query image. A negative sample may include less than ten inlier points. A negative sample may have less than half of the sparse points overlapping with the sparse points of the query image.

[0346] The method 2400 may include computing (Act 2404), for each image set, a loss by comparing the query image with the positive sample image and the negative sample image. The method 2400 may include modifying (Act 2406) the artificial neural network based on the computed loss such that a distance between a frame descriptor generated by the artificial neural network for the query image and a frame descriptor for the positive sample image

is less than a distance between the frame descriptor for the query image and a frame descriptor for the negative sample image.

[0347] It should be appreciated that although methods and apparatus configured to generate global descriptors for individual images are described above, methods and apparatus may be configured to generate descriptors for individual maps. For example, a map may include a plurality of key frames, each of which may have a frame descriptor as described above. A max pool unit may analyze the frame descriptors of the map's key frames and combines the frame descriptors into a unique map descriptor for the map.

[0348] Further, it should be appreciated that other architectures may be used for processing as described above. For example, separate neural networks are described for generating DSF descriptors and frame descriptors. Such an approach is computationally efficient. However, in some embodiments, the frame descriptors may be generated from selected feature points, without first generating DSF descriptors.

Ranking and Merging Maps

[0349] Described herein are methods and apparatus for ranking and merging a plurality of environment maps in a cross reality system. Map merging may enable maps representing overlapping portions of the physical world to be combined to represent a larger area. Ranking maps may enable efficiently performing techniques as described herein, including map merging, that involve selecting a map from a set of maps based on similarity. In some embodiments, for example, a set of canonical maps formatted in a way that they may be accessed by any of a number of XR devices, may be maintained by the system. These canonical maps may be formed by merging selected tracking maps from those devices with other tracking maps or previously stored canonical maps. The canonical maps may be ranked, for example, for use in selecting one or more canonical maps to merge with a new tracking map and/or to select one or more canonical maps from the set to use within a device.

[0350] To provide realistic XR experiences to users, the XR system must know the user's physical surroundings in order to correctly correlate locations of virtual objects in relation to real objects. Information about a user's physical surroundings may be obtained from an environment map for the user's location.

[0351] The inventors have recognized and appreciated that an XR system could provide an enhanced XR experience to multiple users sharing a same world, comprising real and/or virtual content, by enabling efficient sharing of environment maps of the real/physical world collected by multiple users, whether those users are present in the world at the same or different times. However, there are significant challenges in providing such a system. Such a system may store multiple maps generated by multiple users and/or the system may store multiple maps generated at different times. For operations that might be performed with a previously generated map, such as localization, for example as described above, substantial processing may be required to identify a relevant environment map of a same world (e.g. same real world location) from all the environment maps collected in an XR system. In some embodiments, there may only be a small number of environment maps a device could access, for example for localization. In some embodiments, there may be a large number of environment maps a device could

access. The inventors have recognized and appreciated techniques to quickly and accurately rank the relevance of environment maps out of all possible environment maps, such as the universe of all canonical maps **120** in FIG. 28, for example. A high ranking map may then be selected for further processing, such as to render virtual objects on a user display realistically interacting with the physical world around the user or merging map data collected by that user with stored maps to create larger or more accurate maps.

[0352] In some embodiments, a stored map that is relevant to a task for a user at a location in the physical world may be identified by filtering stored maps based on multiple criteria. Those criteria may indicate comparisons of a tracking map, generated by the wearable device of the user in the location, to candidate environment maps stored in a database. The comparisons may be performed based on metadata associated with the maps, such as a Wi-Fi fingerprint detected by the device generating the map and/or set of BSSID's to which the device was connected while forming the map. The comparisons may also be performed based on compressed or uncompressed content of the map. Comparisons based on a compressed representation may be performed, for example, by comparison of vectors computed from map content. Comparisons based on un-compressed maps may be performed, for example, by localizing the tracking map within the stored map, or vice versa. Multiple comparisons may be performed in an order based on computation time needed to reduce the number of candidate maps for consideration, with comparisons involving less computation being performed earlier in the order than other comparisons requiring more computation.

[0353] FIG. 26 depicts an AR system **800** configured to rank and merge one or more environment maps, according to some embodiments. The AR system may include a passable world model **802** of an AR device. Information to populate the passable world model **802** may come from sensors on the AR device, which may include computer executable instructions stored in a processor **804** (e.g., a local data processing module **570** in FIG. 4), which may perform some or all of the processing to convert sensor data into a map. Such a map may be a tracking map, as it can be built as sensor data is collected as the AR device operates in a region. Along with that tracking map, area attributes may be supplied so as to indicate the area that the tracking map represents. These area attributes may be a geographic location identifier, such as coordinates presented as latitude and longitude or an ID used by the AR system to represent a location. Alternatively or additionally, the area attributes may be measured characteristics that have a high likelihood of being unique for that area. The area attributes, for example, may be derived from parameters of wireless networks detected in the area. In some embodiments, the area attribute may be associated with a unique address of an access-point the AR system is nearby and/or connected to. For example, the area attribute may be associated with a MAC address or basic service set identifiers (BSSIDs) of a 5G base station/router, a Wi-Fi router, and the like.

[0354] In the example of FIG. 26, the tracking maps may be merged with other maps of the environment. A map rank portion **806** receives tracking maps from the device PW **802** and communicates with a map database **808** to select and rank environment maps from the map database **808**. Higher ranked, selected maps are sent to a map merge portion **810**.

[0355] The map merge portion **810** may perform merge processing on the maps sent from the map rank portion **806**. Merge processing may entail merging the tracking map with some or all of the ranked maps and transmitting the new, merged maps to a passable world model **812**. The map merge portion may merge maps by identifying maps that depict overlapping portions of the physical world. Those overlapping portions may be aligned such that information in both maps may be aggregated into a final map. Canonical maps may merged with other canonical maps and/or tracking maps.

[0356] The aggregation may entail extending one map with information from another map. Alternatively or additionally, aggregation may entail adjusting the representation of the physical world in one map, based on information in another map. A later map, for example, may reveal that objects giving rise to feature points have moved, such that the map may be updated based on later information. Alternatively, two maps may characterize the same region with different feature points and aggregating may entail selecting a set of feature points from the two maps to better represent that region. Regardless of the specific processing that occurs in the merging process, in some embodiments, PCFs from all maps that are merged may be retained, such that applications positioning content with respect to them may continue to do so. In some embodiments, merging of maps may result in redundant persistent poses, and some of the persistent poses may be deleted. When a PCF is associated with a persistent pose that is to be deleted, merging maps may entail modifying the PCF to be associated with a persistent pose remaining in the map after merging.

[0357] In some embodiments, as maps are extended and or updated, they may be refined. Refinement may entail computation to reduce internal inconsistency between feature points that likely represent the same object in the physical world. Inconsistency may result from inaccuracies in the poses associated with key frames supplying feature points that represent the same objects in the physical world. Such inconsistency may result, for example, from an XR device computing poses relative to a tracking map, which in turn is built based on estimating poses, such that errors in pose estimation accumulate, creating a “drift” in pose accuracy over time. By performing a bundle adjustment or other operation to reduce inconsistencies of the feature points from multiple key frames, the map may be refined.

[0358] Upon a refinement, the location of a persistent point relative to the origin of a map may change. Accordingly, the transformation associated with that persistent point, such as a persistent pose or a PCF, may change. In some embodiments, the XR system, in connection with map refinement (whether as part of a merge operation or performed for other reasons) may re-compute transformations associated with any persistent points that have changed. These transformations might be pushed from a component computing the transformations to a component using the transformation such that any uses of the transformations may be based on the updated location of the persistent points.

[0359] Passable world model **812** may be a cloud model, which may be shared by multiple AR devices. Passable world model **812** may store or otherwise have access to the environment maps in map database **808**. In some embodiments, when a previously computed environment map is updated, the prior version of that map may be deleted so as

to remove out of date maps from the database. In some embodiments, when a previously computed environment map is updated, the prior version of that map may be archived enabling retrieving/viewing prior versions of an environment. In some embodiments, permissions may be set such that only AR systems having certain read/write access may trigger prior versions of maps being deleted/archived.

[0360] These environment maps created from tracking maps supplied by one or more AR devices/systems may be accessed by AR devices in the AR system. The map rank portion 806 also may be used in supplying environment maps to an AR device. The AR device may send a message requesting an environment map for its current location, and map rank portion 806 may be used to select and rank environment maps relevant to the requesting device.

[0361] In some embodiments, the AR system 800 may include a downsample portion 814 configured to receive the merged maps from the cloud PW 812. The received merged maps from the cloud PW 812 may be in a storage format for the cloud, which may include high resolution information, such as a large number of PCFs per square meter or multiple image frames or a large set of feature points associated with a PCF. The downsample portion 814 may be configured to downsample the cloud format maps to a format suitable for storage on AR devices. The device format maps may have less data, such as fewer PCFs or less data stored for each PCF to accommodate the limited local computing power and storage space of AR devices.

[0362] FIG. 27 is a simplified block diagram illustrating a plurality of canonical maps 120 that may be stored in a remote storage medium, for example, a cloud. Each canonical map 120 may include a plurality of canonical map identifiers indicating the canonical map's location within a physical space, such as somewhere on the planet earth. These canonical map identifiers may include one or more of the following identifiers: area identifiers represented by a range of longitudes and latitudes, frame descriptors (e.g., global feature string 316 in FIG. 21), Wi-Fi fingerprints, feature descriptors (e.g., feature descriptors 310 in FIG. 21), and device identities indicating one or more devices that contributed to the map. In the illustrated example, the canonical maps 120 are disposed geographically in a two-dimensional pattern as they may exist on a surface of the earth. The canonical maps 120 may be uniquely identifiable by corresponding longitudes and latitudes because any canonical maps that have overlapping longitudes and latitudes may be merged into a new canonical map.

[0363] FIG. 28 is a schematic diagram illustrating a method of selecting canonical maps, which may be used for localizing a new tracking map to one or more canonical maps, according to some embodiment. The method may start from accessing (Act 120) a universe of canonical maps 120, which may be stored, as an example, in a database in a passable world (e.g., the passable world module 538). The universe of canonical maps may include canonical maps from all previously visited locations. An XR system may filter the universe of all canonical maps to a small subset or just a single map. It should be appreciated that, in some embodiments, it may not be possible to send all the canonical maps to a viewing device due to bandwidth restrictions. Selecting a subset selected as being likely candidates for matching the tracking map to send to the device may reduce bandwidth and latency associated with accessing a remote database of maps.

[0364] The method may include filtering (Act 300) the universe of canonical maps based on areas with predetermined size and shapes. In the illustrated example in FIG. 27, each square may represent an area. Each square may cover 50 m×50 m. Each square may have six neighboring areas. In some embodiments, Act 300 may select at least one matching canonical map 120 covering longitude and latitude that include that longitude and latitude of the position identifier received from an XR device, as long as at least one map exists at that longitude and latitude. In some embodiments, the Act 300 may select at least one neighboring canonical map covering longitude and latitude that are adjacent the matching canonical map. In some embodiments, the Act 300 may select a plurality of matching canonical maps and a plurality of neighboring canonical maps. The Act 300 may, for example, reduce the number of canonical maps approximately ten times, for example, from thousands to hundreds to form a first filtered selection. Alternatively or additionally, criteria other than latitude and longitude may be used to identify neighboring maps. An XR device, for example, may have previously localized with a canonical map in the set as part of the same session. A cloud service may retain information about the XR device, including maps previously localized to. In this example, the maps selected at Act 300 may include those that cover an area adjacent to the map to which the XR device localized to.

[0365] The method may include filtering (Act 302) the first filtered selection of canonical maps based on Wi-Fi fingerprints. The Act 302 may determine a latitude and longitude based on a Wi-Fi fingerprint received as part of the position identifier from an XR device. The Act 302 may compare the latitude and longitude from the Wi-Fi fingerprint with latitude and longitude of the canonical maps 120 to determine one or more canonical maps that form a second filtered selection. The Act 302 may reduce the number of canonical maps approximately ten times, for example, from hundreds to tens of canonical maps (e.g., 50) that form a second selection. For example, a first filtered selection may include 130 canonical maps and the second filtered selection may include 50 of the 130 canonical maps and may not include the other 80 of the 130 canonical maps.

[0366] The method may include filtering (Act 304) the second filtered selection of canonical maps based on key frames. The Act 304 may compare data representing an image captured by an XR device with data representing the canonical maps 120. In some embodiments, the data representing the image and/or maps may include feature descriptors (e.g., DSF descriptors in FIG. 25) and/or global feature strings (e.g., 316 in FIG. 21). The Act 304 may provide a third filtered selection of canonical maps. In some embodiments, the output of Act 304 may only be five of the 50 canonical maps identified following the second filtered selection, for example. The map transmitter 122 then transmits the one or more canonical maps based on the third filtered selection to the viewing device. The Act 304 may reduce the number of canonical maps for approximately ten times, for example, from tens to single digits of canonical maps (e.g., 5) that form a third selection. In some embodiments, an XR device may receive canonical maps in the third filtered selection, and attempt to localize into the received canonical maps.

[0367] For example, the Act 304 may filter the canonical maps 120 based on the global feature strings 316 of the canonical maps 120 and the global feature string 316 that is

based on an image that is captured by the viewing device (e.g. an image that may be part of the local tracking map for a user). Each one of the canonical maps **120** in FIG. 27 thus has one or more global feature strings **316** associated therewith. In some embodiments, the global feature strings **316** may be acquired when an XR device submits images or feature details to the cloud and the cloud processes the image or feature details to generate global feature strings **316** for the canonical maps **120**.

[0368] In some embodiments, the cloud may receive feature details of a live/new/current image captured by a viewing device, and the cloud may generate a global feature string **316** for the live image. The cloud may then filter the canonical maps **120** based on the live global feature string **316**. In some embodiments, the global feature string may be generated on the local viewing device. In some embodiments, the global feature string may be generated remotely, for example on the cloud. In some embodiments, a cloud may transmit the filtered canonical maps to an XR device together with the global feature strings **316** associated with the filtered canonical maps. In some embodiments, when the viewing device localizes its tracking map to the canonical map, it may do so by matching the global feature strings **316** of the local tracking map with the global feature strings of the canonical map.

[0369] It should be appreciated that an operation of an XR device may not perform all of the Acts (300, 302, 304). For example, if a universe of canonical maps are relatively small (e.g., 500 maps), an XR device attempting to localize may filter the universe of canonical maps based on Wi-Fi fingerprints (e.g., Act 302) and Key Frame (e.g., Act 304), but omit filtering based on areas (e.g., Act 300). Moreover, it is not necessary that maps in their entirities be compared. In some embodiments, for example, a comparison of two maps may result in identifying common persistent points, such as persistent poses or PCFs that appear in both the new map the selected map from the universe of maps. In that case, descriptors may be associated with persistent points, and those descriptors may be compared.

[0370] FIG. 29 is flow chart illustrating a method **900** of selecting one or more ranked environment maps, according to some embodiments. In the illustrated embodiment, the ranking is performed for a user's AR device that is creating a tracking map. Accordingly, the tracking map is available for use in ranking environment maps. In embodiments in which the tracking map is not available, some or all of portions of the selection and ranking of environment maps that do not expressly rely on the tracking map may be used.

[0371] The method **900** may start at Act **902**, where a set of maps from a database of environment maps (which may be formatted as canonical maps) that are in the neighborhood of the location where the tracking map was formed may be accessed and then filtered for ranking. Additionally, at Act **902**, at least one area attribute for the area in which the user's AR device is operating is determined. In scenarios in which the user's AR device is constructing a tracking map, the area attributes may correspond to the area over which the tracking map was created. As a specific example, the area attributes may be computed based on received signals from access points to computer networks while the AR device was computing the tracking map.

[0372] FIG. 30 depicts an exemplary map rank portion **806** of the AR system **800**, according to some embodiments. The map rank portion **806** may be executing in a cloud

computing environment, as it may include portions executing on AR devices and portions executing on a remote computing system such as a cloud. The map rank portion **806** may be configured to perform at least a portion of the method **900**.

[0373] FIG. 31A depicts an example of area attributes AA1-AA8 of a tracking map (TM) **1102** and environment maps CM1-CM4 in a database, according to some embodiments. As illustrated, an environment map may be associated to multiple area attributes. The area attributes AA1-AA8 may include parameters of wireless networks detected by the AR device computing the tracking map **1102**, for example, basic service set identifiers (BSSIDs) of networks to which the AR device are connected and/or the strength of the received signals of the access points to the wireless networks through, for example, a network tower **1104**. The parameters of the wireless networks may comply with protocols including Wi-Fi and 5G NR. In the example illustrated in FIG. 32, the area attributes are a fingerprint of the area in which the user AR device collected sensor data to form the tracking map.

[0374] FIG. 31B depicts an example of the determined geographic location **1106** of the tracking map **1102**, according to some embodiments. In the illustrated example, the determined geographic location **1106** includes a centroid point **1110** and an area **1108** circling around the centroid point. It should be appreciated that the determination of a geographic location of the present application is not limited to the illustrated format. A determined geographic location may have any suitable formats including, for example, different area shapes. In this example, the geographic location is determined from area attributes using a database relating area attributes to geographic locations. Databases are commercially available, for example, databases that relate Wi-Fi fingerprints to locations expressed as latitude and longitude and may be used for this operation.

[0375] In the embodiment of FIG. 29, a map database, containing environment maps may also include location data for those maps, including latitude and longitude covered by the maps. Processing at Act **902** may entail selecting from that database a set of environment maps that covers the same latitude and longitude determined for the area attributes of the tracking map.

[0376] Act **904** is a first filtering of the set of environment maps accessed in Act **902**. In Act **902**, environment maps are retained in the set based on proximity to the geolocation of the tracking map. This filtering step may be performed by comparing the latitude and longitude associated with the tracking map and the environment maps in the set.

[0377] FIG. 32 depicts an example of Act **904**, according to some embodiments. Each area attribute may have a corresponding geographic location **1202**. The set of environment maps may include the environment maps with at least one area attribute that has a geographic location overlapping with the determined geographic location of the tracking map. In the illustrated example, the set of identified environment maps includes environment maps CM1, CM2, and CM4, each of which has at least one area attribute that has a geographic location overlapping with the determined geographic location of the tracking map **1102**. The environment map CM3 associated with the area attribute AA6 is not included in the set because it is outside the determined geographic location of the tracking map.

[0378] Other filtering steps may also be performed on the set of environment maps to reduce/rank the number of environment maps in the set that is ultimately processed (such as for map merge or to provide passable world information to a user device). The method 900 may include filtering (Act 906) the set of environment maps based on similarity of one or more identifiers of network access points associated with the tracking map and the environment maps of the set of environment maps. During the formation of a map, a device collecting sensor data to generate the map may be connected to a network through a network access point, such as through Wi-Fi or similar wireless communication protocol. The access points may be identified by BSSID. The user device may connect to multiple different access points as it moves through an area collecting data to form a map. Likewise, when multiple devices supply information to form a map, the devices may have connected through different access points, so there may be multiple access points used in forming the map for that reason too. Accordingly, there may be multiple access points associated with a map, and the set of access points may be an indication of location of the map. Strength of signals from an access point, which may be reflected as an RSSI value, may provide further geographic information. In some embodiments, a list of BSSID and RSSI values may form the area attribute for a map.

[0379] In some embodiments, filtering the set of environment maps based on similarity of the one or more identifiers of the network access points may include retaining in the set of environment maps environment maps with the highest Jaccard similarity to the at least one area attribute of the tracking map based on the one or more identifiers of network access points. FIG. 33 depicts an example of Act 906, according to some embodiments. In the illustrated example, a network identifier associated with the area attribute AA7 may be determined as the identifier for the tracking map 1102. The set of environment maps after Act 906 includes environment map CM2, which may have area attributes within higher Jaccard similarity to AA7, and environment map CM4, which also include the area attributes AA7. The environment map CM1 is not included in the set because it has the lowest Jaccard similarity to AA7.

[0380] Processing at Acts 902-906 may be performed based on metadata associated with maps and without actually accessing the content of the maps stored in a map database. Other processing may involve accessing the content of the maps. Act 908 indicates accessing the environment maps remaining in the subset after filtering based on metadata. It should be appreciated that this act may be performed either earlier or later in the process, if subsequent operations can be performed with accessed content.

[0381] The method 900 may include filtering (Act 910) the set of environment maps based on similarity of metrics representing content of the tracking map and the environment maps of the set of environment maps. The metrics representing content of the tracking map and the environment maps may include vectors of values computed from the contents of the maps. For example, the Deep Key Frame descriptor, as described above, computed for one or more key frames used in forming a map may provide a metric for comparison of maps, or portions of maps. The metrics may be computed from the maps retrieved at Act 908 or may be pre-computed and stored as metadata associated with those maps. In some embodiments, filtering the set of environment

maps based on similarity of metrics representing content of the tracking map and the environment maps of the set of environment maps, may include retaining in the set of environment maps environment maps with the smallest vector distance between a vector of characteristics of the tracking map and vectors representing environment maps in the set of environment maps.

[0382] The method 900 may include further filtering (Act 912) the set of environment maps based on degree of match between a portion of the tracking map and portions of the environment maps of the set of environment maps. The degree of match may be determined as a part of a localization process. As a non-limiting example, localization may be performed by identifying critical points in the tracking map and the environment map that are sufficiently similar as they could represent the same portion of the physical world. In some embodiments, the critical points may be features, feature descriptors, key frames/key rigs, persistent poses, and/or PCFs. The set of critical points in the tracking map might then be aligned to produce a best fit with the set of critical points in the environment map. A mean square distance between the corresponding critical points might be computed and, if below a threshold for a particular region of the tracking map, used as an indication that the tracking map and the environment map represent the same region of the physical world.

[0383] In some embodiments, filtering the set of environment maps based on degree of match between a portion of the tracking map and portions of the environment maps of the set of environment maps may include computing a volume of a physical world represented by the tracking map that is also represented in an environment map of the set of environment maps, and retaining in the set of environment maps environment maps with a larger computed volume than environment maps filtered out of the set. FIG. 34 depicts an example of Act 912, according to some embodiments. In the illustrated example, the set of environment maps after Act 912 includes environment map CM4, which has an area 1402 matched with an area of the tracking map 1102. The environment map CM1 is not included in the set because it has no area matched with an area of the tracking map 1102.

[0384] In some embodiments, the set of environment maps may be filtered in the order of Act 906, Act 910, and Act 912. In some embodiments, the set of environment maps may be filtered based on Act 906, Act 910, and Act 912, which may be performed in an order based on processing required to perform the filtering, from lowest to highest. The method 900 may include loading (Act 914) the set of environment maps and data.

[0385] In the illustrated example, a user database stores area identities indicating areas that AR devices were used in. The area identities may be area attributes, which may include parameters of wireless networks detected by the AR devices when in use. A map database may store multiple environment maps constructed from data supplied by the AR devices and associated metadata. The associated metadata may include area identities derived from the area identities of the AR devices that supplied data from which the environment maps were constructed. An AR device may send a message to a PW module indicating a new tracking map is created or being created. The PW module may compute area identifiers for the AR device and updates the user database based on the received parameters and/or the computed area

identifiers. The PW module may also determine area identifiers associated with the AR device requesting the environment maps, identify sets of environment maps from the map database based on the area identifiers, filter the sets of environment maps, and transmit the filtered sets of environment maps to the AR devices. In some embodiments, the PW module may filter the sets of environment maps based on one or more criteria including, for example, a geographic location of the tracking map, similarity of one or more identifiers of network access points associated with the tracking map and the environment maps of the set of environment maps, similarity of metrics representing contents of the tracking map and the environment maps of the set of environment maps, and degree of match between a portion of the tracking map and portions of the environment maps of the set of environment maps.

[0386] Having thus described several aspects of some embodiments, it is to be appreciated that various alterations, modifications, and improvements will readily occur to those skilled in the art. As one example, embodiments are described in connection with an augmented (AR) environment. It should be appreciated that some or all of the techniques described herein may be applied in an MR environment or more generally in other XR environments, and in VR environments.

[0387] As another example, embodiments are described in connection with devices, such as wearable devices. It should be appreciated that some or all of the techniques described herein may be implemented via networks (such as cloud), discrete applications, and/or any suitable combinations of devices, networks, and discrete applications.

[0388] Further, FIG. 29 provides examples of criteria that may be used to filter candidate maps to yield a set of high ranking maps. Other criteria may be used instead of or in addition to the described criteria. For example, if multiple candidate maps have similar values of a metric used for filtering out less desirable maps, characteristics of the candidate maps may be used to determine which maps are retained as candidate maps or filtered out. For example, larger or more dense candidate maps may be prioritized over smaller candidate maps. In some embodiments, FIGS. 27-28 may describe all or part of the systems and methods described in FIGS. 29-34.

[0389] FIGS. 35 and 36 are schematic diagrams illustrating an XR system configured to rank and merge a plurality of environment maps, according to some embodiments. In some embodiments, a passable world (PW) may determine when to trigger ranking and/or merging the maps. In some embodiments, determining a map to be used may be based at least partly on deep key frames described above in relation to FIGS. 21-25, according to some embodiments.

[0390] FIG. 37 is a block diagram illustrating a method 3700 of creating environment maps of a physical world, according to some embodiments. The method 3700 may start from localizing (Act 3702) a tracking map captured by an XR device worn by a user to a group of canonical maps (e.g., canonical maps selected by the method of FIG. 28 and/or the method 900 of FIG. 29). The Act 3702 may include localizing keyrigs of the tracking map into the group of canonical maps. The localization result of each keyrig may include the keyrig's localized pose and a set of 2D-to-3D feature correspondences.

[0391] In some embodiments, the method 3700 may include splitting (Act 3704) a tracking map into connected

components, which may enable merging maps robustly by merging connected pieces. Each connected component may include keyrigs that are within a predetermined distance. The method 3700 may include merging (Act 3706) the connected components that are larger than a predetermined threshold into one or more canonical maps, and removing the merged connected components from the tracking map. [0392] In some embodiments, the method 3700 may include merging (Act 3708) canonical maps of the group that are merged with the same connected components of the tracking map. In some embodiments, the method 3700 may include promoting (Act 3710) the remaining connected components of the tracking map that has not been merged with any canonical maps to be a canonical map. In some embodiments, the method 3700 may include merging (Act 3712) persistent poses and/or PCFs of the tracking maps and the canonical maps that are merged with at least one connected component of the tracking map. In some embodiments, the method 3700 may include finalizing (Act 3714) the canonical maps by, for example, fusing map points and pruning redundant keyrigs.

[0393] FIGS. 38A and 38B illustrate an environment map 3800 created by updating a canonical map 700, which may be promoted from the tracking map 700 (FIG. 7) with a new tracking map, according to some embodiments. As illustrated and described with respect to FIG. 7, the canonical map 700 may provide a floor plan 706 of reconstructed physical objects in a corresponding physical world, represented by points 702. In some embodiments, a map point 702 may represent a feature of a physical object that may include multiple features. A new tracking map may be captured about the physical world and uploaded to a cloud to merge with the map 700. The new tracking map may include map points 3802, and keyrigs 3804, 3806. In the illustrated example, keyrigs 3804 represent keyrigs that are successfully localized to the canonical map by, for example, establishing a correspondence with a keyrig 704 of the map 700 (as illustrated in FIG. 38B). On the other hand, keyrigs 3806 represent keyrigs that have not been localized to the map 700. Keyrigs 3806 may be promoted to a separate canonical map in some embodiments.

[0394] FIGS. 39A to 39F are schematic diagrams illustrating an example of a cloud-based persistent coordinate system providing a shared experience for users in the same physical space. FIG. 39A shows that a canonical map 4814, for example, from a cloud, is received by the XR devices worn by the users 4802A and 4802B of FIGS. 20A-20C. The canonical map 4814 may have a canonical coordinate frame 4806C. The canonical map 4814 may have a PCF 4810C with a plurality of associated PPs (e.g., 4818A, 4818B in FIG. 39C).

[0395] FIG. 39B shows that the XR devices established relationships between their respective world coordinate system 4806A, 4806B with the canonical coordinate frame 4806C. This may be done, for example, by localizing to the canonical map 4814 on the respective devices. Localizing the tracking map to the canonical map may result, for each device, a transformation between its local world coordinate system and the coordinate system of the canonical map.

[0396] FIG. 39C shows that, as a result of localization, a transformation can be computed (e.g., transformation 4816A, 4816B) between a local PCF (e.g., PCFs 4810A, 4810B) on the respective device to a respective persistent pose (e.g., PPs 4818A, 4818B) on the canonical map. With

these transformations, each device may use its local PCFs, which can be detected locally on the device by processing images detected with sensors on the device, to determine where with respect to the local device to display virtual content attached to the PPs **4818A**, **4818B** or other persistent points of the canonical map. Such an approach may accurately position virtual content with respect to each user and may enable each user to have the same experience of the virtual content in the physical space.

[0397] FIG. 39D shows a persistent pose snapshot from the canonical map to the local tracking maps. As can be seen, the local tracking maps are connected to one another via the persistent poses. FIG. 39E shows that the PCF **4810A** on the device worn by the user **4802A** is accessible in the device worn by the user **4802B** through PPs **4818A**. FIG. 39F shows that the tracking maps **4804A**, **4804B** and the canonical **4814** may merge. In some embodiments, some PCFs may be removed as a result of merging. In the illustrated example, the merged map includes the PCF **4810C** of the canonical map **4814** but not the PCFs **4810A**, **4810B** of the tracking maps **4804A**, **4804B**. The PPs previously associated with the PCFs **4810A**, **4810B** may be associated with the PCF **4810C** after the maps merge.

EXAMPLES

[0398] FIGS. 40 and 41 illustrate an example of using a tracking map by the first XR device **12.1** of FIG. 9. FIG. 40 is a two-dimensional representation of a three-dimensional first local tracking map (Map 1), which may be generated by the first XR device of FIG. 9, according to some embodiments. FIG. 41 is a block diagram illustrating uploading Map 1 from the first XR device to the server of FIG. 9, according to some embodiments.

[0399] FIG. 40 illustrates Map 1 and virtual content (Content**123** and Content**456**) on the first XR device **12.1**. Map 1 has an origin (Origin 1). Map 1 includes a number of PCFs (PCF a to PCF d). From the perspective of the first XR device **12.1**, PCF a, by way of example, is located at the origin of Map 1 and has X, Y, and Z coordinates of (0,0,0) and PCF b has X, Y, and Z coordinates (-1,0,0). Content**123** is associated with PCF a. In the present example, Content**123** has an X, Y, and Z relationship relative to PCF a of (1,0,0). Content**456** has a relationship relative to PCF b. In the present example, Content**456** has an X, Y, and Z relationship of (1,0,0) relative to PCF b.

[0400] In FIG. 41, the first XR device **12.1** uploads Map 1 to the server **20**. In this example, as the server stores no canonical map for the same region of the physical world represented by the tracking map, and the tracking map is stored as an initial canonical map. The server **20** now has a canonical map based on Map 1. The first XR device **12.1** has a canonical map that is empty at this stage. The server **20**, for purposes of discussion, and in some embodiments, includes no other maps other than Map 1. No maps are stored on the second XR device **12.2**.

[0401] The first XR device **12.1** also transmits its Wi-Fi signature data to the server **20**. The server **20** may use the Wi-Fi signature data to determine a rough location of the first XR device **12.1** based on intelligence gathered from other devices that have, in the past, connected to the server **20** or other servers together with the GPS locations of such other devices that have been recorded. The first XR device **12.1** may now end the first session (See FIG. 8) and may disconnect from the server **20**.

[0402] FIG. 42 is a schematic diagram illustrating the XR system of FIG. 16, showing the second user **14.2** has initiated a second session using a second XR device of the XR system after the first user **14.1** has terminated a first session, according to some embodiments. FIG. 43A is a block diagram showing the initiation of a second session by a second user **14.2**. The first user **14.1** is shown in phantom lines because the first session by the first user **14.1** has ended. The second XR device **12.2** begins to record objects. Various systems with varying degrees of granulation may be used by the server **20** to determine that the second session by the second XR device **12.2** is in the same vicinity of the first session by the first XR device **12.1**. For example, Wi-Fi signature data, global positioning system (GPS) positioning data, GPS data based on Wi-Fi signature data, or any other data that indicates a location may be included in the first and second XR devices **12.1** and **12.2** to record their locations. Alternatively, the PCFs that are identified by the second XR device **12.2** may show a similarity to the PCFs of Map 1.

[0403] As shown in FIG. 43B, the second XR device boots up and begins to collect data, such as images **1110** from one or more cameras **44**, **46**. As shown in FIG. 14, in some embodiments, an XR device (e.g. the second XR device **12.2**) may collect one or more images **1110** and perform image processing to extract one or more features/interest points **1120**. Each feature may be converted to a descriptor **1130**. In some embodiments, the descriptors **1130** may be used to describe a key frame **1140**, which may have the position and direction of the associated image attached. One or more key frames **1140** may correspond to a single persistent pose **1150**, which may be automatically generated after a threshold distance from the previous persistent pose **1150**, e.g., 3 meters. One or more persistent poses **1150** may correspond to a single PCF **1160**, which may be automatically generated after a pre-determined distance, e.g. every 5 meters. Over time as the user continues to move around the user's environment, and the XR device continues to collect more data, such as images **1110**, additional PCFs (e.g., PCF **3** and PCF **4**, **5**) may be created. One or more applications **1180** may run on the XR device and provide virtual content **1170** to the XR device for presentation to the user. The virtual content may have an associated content coordinate frame which may be placed relative to one or more PCFs. As shown in FIG. 43B, the second XR device **12.2** creates three PCFs. In some embodiments, the second XR device **12.2** may try to localize into one or more canonical maps stored on the server **20**.

[0404] In some embodiments, as shown in FIG. 43C, the second XR device **12.2** may download the canonical map **120** from the server **20**. Map 1 on the second XR device **12.2** includes PCFs a to d and Origin 1. In some embodiments, the server **20** may have multiple canonical maps for various locations and may determine that the second XR device **12.2** is in the same vicinity as the vicinity of the first XR device **12.1** during the first session and sends the second XR device **12.2** the canonical map for that vicinity.

[0405] FIG. 44 shows the second XR device **12.2** beginning to identify PCFs for purposes of generating Map 2. The second XR device **12.2** has only identified a single PCF, namely PCF **1.2**. The X, Y, and Z coordinates of PCF **1.2** for the second XR device **12.2** may be (1,1,1). Map 2 has its own origin (Origin 2), which may be based on the headpose of device **2** at device start-up for the current headpose session. In some embodiments, the second XR device **12.2**

may immediately attempt to localize Map 2 to the canonical map. In some embodiments, Map 2 may not be able to localize into Canonical Map (Map 1) (i.e. localization may fail) because the system does not recognize any or enough overlap between the two maps. Localization may be performed by identifying a portion of the physical world represented in a first map that is also represented in a second map, and computing a transformation between the first map and the second map required to align those portions. In some embodiments, the system may localize based on PCF comparison between the local and canonical maps. In some embodiments, the system may localize based on persistent pose comparison between the local and canonical maps. In some embodiments, the system may localize based on key frame comparison between the local and canonical maps.

[0406] FIG. 45 shows Map 2 after the second XR device 12.2 has identified further PCFs (PCF 1,2, PCF 3, PCF 4,5) of Map 2. The second XR device 12.2 again attempts to localize Map 2 to the canonical map. Because Map 2 has expanded to overlap with at least a portion of the Canonical Map, the localization attempt will succeed. In some embodiments, the overlap between the local tracking map, Map 2, and the Canonical Map may be represented by PCFs, persistent poses, key frames, or any other suitable intermediate or derivative construct.

[0407] Furthermore, the second XR device 12.2 has associated Content123 and Content456 to PCFs 1,2 and PCF 3 of Map 2. Content123 has X, Y, and Z coordinates relative to PCF 1,2 of (1,0,0). Similarly, the X, Y, and Z coordinates of Content456 relative to PCF 3 in Map 2 are (1,0,0).

[0408] FIGS. 46A and 46B illustrate a successful localization of Map 2 to the canonical map. Localization may be based on matching features in one map to the other. With an appropriate transformation, here involving both translation and rotation of one map with respect to the other, the overlapping area/volume/section of the maps 1410 represent the common parts to Map 1 and the canonical map. Since Map 2 created PCFs 3 and 4,5 before localizing, and the Canonical map created PCFs a and c before Map 2 was created, different PCFs were created to represent the same volume in real space (e.g., in different maps).

[0409] As shown in FIG. 47, the second XR device 12.2 expands Map 2 to include PCFs a-d from the Canonical Map. The inclusion of PCFs a-d represents the localization of Map 2 to the Canonical Map. In some embodiments, the XR system may perform an optimization step to remove duplicate PCFs from overlapping areas, such as the PCFs in 1410, PCF 3 and PCF 4,5. After Map 2 localizes, the placement of virtual content, such as Content456 and Content123 will be relative to the closest updated PCFs in the updated Map 2. The virtual content appears in the same real-world location relative to the user, despite the changed PCF attachment for the content, and despite the updated PCFs for Map 2.

[0410] As shown in FIG. 48, the second XR device 12.2 continues to expand Map 2 as further PCFs (e.g., PCFs e, f, g, and h) are identified by the second XR device 12.2, for example as the user walks around the real world. It can also be noted that Map 1 has not expanded in FIGS. 47 and 48.

[0411] Referring to FIG. 49, the second XR device 12.2 uploads Map 2 to the server 20. The server 20 stores Map 2 together with the canonical map. In some embodiments, Map 2 may upload to the server 20 when the session ends for the second XR device 12.2.

[0412] The canonical map within the server 20 now includes PCF i which is not included in Map 1 on the first XR device 12.1. The canonical map on the server 20 may have expanded to include PCF i when a third XR device (not shown) uploaded a map to the server 20 and such a map included PCF i.

[0413] In FIG. 50, the server 20 merges Map 2 with the canonical map to form a new canonical map. The server 20 determines that PCFs a to d are common to the canonical map and Map 2. The server expands the canonical map to include PCFs e to h and PCF 1,2 from Map 2 to form a new canonical map. The canonical maps on the first and second XR devices 12.1 and 12.2 are based on Map 1 and are outdated.

[0414] In FIG. 51, the server 20 transmits the new canonical map to the first and second XR devices 12.1 and 12.2. In some embodiments, this may occur when the first XR device 12.1 and second device 12.2 try to localize during a different or new or subsequent session. The first and second XR devices 12.1 and 12.2 proceed as described above to localize their respective local maps (Map 1 and Map 2 respectively) to the new canonical map.

[0415] As shown in FIG. 52, the head coordinate frame 96 or “headpose” is related to the PCFs in Map 2. In some embodiments, the origin of the map, Origin 2, is based on the headpose of second XR device 12.2 at the start of the session. As PCFs are created during the session, the PCFs are placed relative to the world coordinate frame, Origin 2. The PCFs of Map 2 serve as a persistent coordinate frames relative to a canonical coordinate frame, where the world coordinate frame may be a previous session’s world coordinate frame (e.g. Map 1’s Origin 1 in FIG. 40). These coordinate frames are related by the same transformation used to localize Map 2 to the canonical map, as discussed above in connection with FIG. 46B.

[0416] The transformation from the world coordinate frame to the head coordinate frame 96 has been previously discussed with reference to FIG. 9. The head coordinate frame 96 shown in FIG. 52 only has two orthogonal axes that are in a particular coordinate position relative to the PCFs of Map 2, and at particular angles relative to Map 2. It should however be understood that the head coordinate frame 96 is in a three-dimensional location relative to the PCFs of Map 2 and has three orthogonal axes within three-dimensional space.

[0417] In FIG. 53, the head coordinate frame 96 has moved relative to the PCFs of Map 2. The head coordinate frame 96 has moved because the second user 14.2 has moved their head. The user can move their head in six degrees of freedom (6dof). The head coordinate frame 96 can thus move in 6dof, namely in three-dimensions from its previous location in FIG. 52 and about three orthogonal axes relative to the PCFs of Map 2. The head coordinate frame 96 is adjusted when the real object detection camera 44 and inertial measurement unit 48 in FIG. 9 respectively detect real objects and motion of the head unit 22. More information regarding headpose tracking is disclosed in U.S. patent application Ser. No. 16/221,065 entitled “Enhanced Pose Determination for Display Device” and is hereby incorporated by reference in its entirety.

[0418] FIG. 54 shows that sound may be associated with one or more PCFs. A user may, for example, wear headphones or earphones with stereoscopic sound. The location of sound through headphones can be simulated using con-

ventional techniques. The location of sound may be located in a stationary position so that, when the user rotates their head to the left, the location of sound rotates to the right so that the user perceives the sound coming from the same location in the real world. In the present example, location of sound is represented by Sound123 and Sound456. For purposes of discussion, FIG. 54 is similar to FIG. 48 in its analysis. When the first and second users 14.1 and 14.2 are located in the same room at the same or different times, they perceive Sound123 and Sound456 coming from the same locations within the real world.

[0419] FIGS. 55 and 56 illustrate a further implementation of the technology described above. The first user 14.1 has initiated a first session as described with reference to FIG. 8. As shown in FIG. 55, the first user 14.1 has terminated the first session as indicated by the phantom lines. At the end of the first session, the first XR device 12.1 uploaded Map 1 to the server 20. The first user 14.1 has now initiated a second session at a later time than the first session. The first XR device 12.1 does not download Map 1 from the server 20 because Map 1 is already stored on the first XR device 12.1. If Map 1 is lost, then the first XR device 12.1 downloads Map 1 from the server 20. The first XR device 12.1 then proceeds to build PCFs for Map 2, localizes to Map 1, and further develops a canonical map as described above. Map 2 of the first XR device 12.1 is then used for relating local content, a head coordinate frame, local sound, etc. as described above.

[0420] Referring to FIGS. 57 and 58, it may also be possible that more than one user interacts with the server in the same session. In the present example, the first user 14.1 and the second user 14.2 are joined by a third user 14.3 with a third XR device 12.3. Each XR device 12.1, 12.2, and 12.3 begins to generate its own map, namely Map 1, Map 2, and Map 3, respectively. As the XR devices 12.1, 12.2, and 12.3 continue to develop Maps 1, 2, and 3, the maps are incrementally uploaded to the server 20. The server 20 merges Maps 1, 2, and 3 to form a canonical map. The canonical map is then transmitted from the server 20 to each one of the XR devices 12.1, 12.2 and 12.3.

[0421] FIG. 59 illustrates aspects of a viewing method to recover and/or reset headpose, according to some embodiments. In the illustrated example, at Act 1400, the viewing device is powered on. At Act 1410, in response to being powered on, a new session is initiated. In some embodiments, a new session may include establishing headpose. One or more capture devices on a head-mounted frame secured to a head of a user capture surfaces of an environment by first capturing images of the environment and then determining the surfaces from the images. In some embodiments, surface data may be combined with a data from a gravitational sensor to establish headpose. Other suitable methods of establishing headpose may be used.

[0422] At Act 1420, a processor of the viewing device enters a routine for tracking of headpose. The capture devices continue to capture surfaces of the environment as the user moves their head to determine an orientation of the head-mounted frame relative to the surfaces.

[0423] At Act 1430, the processor determines whether headpose has been lost. Headpose may become lost due to “edge” cases, such as too many reflective surfaces, low light, blank walls, being outdoor, etc. that may result in low feature acquisition, or because of dynamic cases such as a crowd that moves and forms part of the map. The routine at

1430 allows for a certain amount of time, for example 10 seconds, to pass to allow enough time to determine whether headpose has been lost. If headpose has not been lost, then the processor returns to 1420 and again enters tracking of headpose.

[0424] If headpose has been lost at Act 1430, the processor enters a routine at 1440 to recover headpose. If headpose is lost due to low light, then a message such as the following message is displayed to the user through a display of the viewing device:

[0425] THE SYSTEM IS DETECTING A LOW LIGHT CONDITION. PLEASE MOVE TO AN AREA WHERE THERE IS MORE LIGHT.

[0426] The system will continue to monitor whether there is sufficient light available and whether headpose can be recovered. The system may alternatively determine that low texture of surfaces is causing headpose to be lost, in which case the user is given the following prompt in the display as a suggestion to improve capturing of surfaces:

[0427] THE SYSTEM CANNOT DETECT ENOUGH SURFACES WITH FINE TEXTURES. PLEASE MOVE TO AN AREA WHERE THE SURFACES ARE LESS ROUGH IN TEXTURE AND MORE REFINED IN TEXTURE.

[0428] At Act 1450, the processor enters a routine to determine whether headpose recovery has failed. If headpose recovery has not failed (i.e. headpose recovery has succeeded), then the processor returns to Act 1420 by again entering tracking of headpose. If headpose recovery has failed, the processor returns to Act 1410 to establish a new session. As part of the new session, all cached data is invalidated, whereafter headpose is established anew. Any suitable method of head tracking may be used in combination with the process described in FIG. 59. U.S. patent application Ser. No. 16/221,065 describes head tracking and is hereby incorporated by reference in its entirety.

Remote Localization

[0429] Various embodiments may utilize remote resources to facilitate persistent and consistent cross reality experiences between individual and/or groups of users. The inventors have recognized and appreciated that the benefits of operation of an XR device with canonical maps as described herein can be achieved without downloading a set of canonical maps, such as is illustrated in FIG. 30. The benefit, for example, may be achieved by sending feature and pose information to a remote service that maintains a set of canonical maps. A device seeking to use a canonical map to position virtual content in locations specified relative to the canonical map may receive from the remote service one or more transformations between the features and the canonical maps. Those transformations may be used on the device, which maintains information about the positions of those features in the physical world, to position virtual content in locations specified with respect to canonical map or to otherwise identify locations in the physical world that are specified with respect to the canonical map.

[0430] In some embodiments, spatial information is captured by an XR device and communicated to a remote service, such as a cloud based service, which uses the spatial information to localize the XR device to a canonical map used by applications or other components of an XR system to specify the location of virtual content with respect to the physical world. Once localized, transforms that link a track-

ing map maintained by the device to the canonical map can be communicated to the device. The transforms may be used, in conjunction with the tracking map, to determine a position in which to render virtual content specified with respect to the canonical map, or otherwise identify locations in the physical world that are specified with respect to the canonical map.

[0431] The inventors have realized that the data needed to be exchanged between a device and a remote localization service can be quite small relative to communicating map data, as might occur when a device communicates a tracking map to a remote service and receives from that service a set of canonical maps for device based localization). In some embodiments, performing localization functions on cloud resources requires only small amount of information to be transmitted from the device to the remote service. It is not a requirement, for example, that a full tracking map be communicated to the remote service to perform localization. In some embodiments, features and pose information, such as might be stored in connection with a persistent pose, as described above, might be transmitted to the remote server. In embodiments in which features are represented by descriptors, as described above, the information uploaded may be even smaller.

[0432] The results returned to the device from the localization service may be one or more transformations that relate the uploaded features to portions of a matching canonical map. Those transformations may be used within the XR system, in conjunction with its tracking map, for identifying locations of virtual content or otherwise identifying locations in the physical world. In embodiments in which persistent spatial information, such as PCFs as described above, are used to specify locations with respect to a canonical map, the localization service may download to the device transformations between the features and one or more PCFs after a successful localization.

[0433] As a result, network bandwidth consumed by communications between an XR device and a remote service for performing localization may be low. The system may therefore support frequent localization, enabling each device interacting with the system to quickly obtain information for positioning virtual content or performing other location-based functions. As a device moves within the physical environment, it may repeat requests for updated localization information. Additionally, a device may frequently obtain updates to the localization information, such as when the canonical maps change, such as through merging of additional tracking maps to expand the map or increase their accuracy.

[0434] Further, uploading features and downloading transformations can enhance privacy in an XR system that shares map information among multiple users by increasing the difficulty of obtaining maps by spoofing. An unauthorized user, for example, may be thwarted from obtaining a map from the system by sending a fake request for a canonical map representing a portion of the physical world in which that unauthorized user is not located. An unauthorized user would be unlikely to have access to the features in the region of the physical world for which it is requesting map information if not physically present in that region. In embodiments in which feature information is formatted as feature descriptions, the difficulty in spoofing feature information in a request for map information would be compounded. Further, when the system returns a transformation intended to be

applied to a tracking map of a device operating in the region about which location information is requested, the information returned by the system is likely to be of little or no use to an imposter.

[0435] According to one embodiment, a localization service is implemented as a cloud based micro-service. In some examples, implementing a cloud-based localization service can help save device compute resources and may enable computations required for localization to be performed with very low latency. Those operations can be supported by nearly infinite compute power or other computing resources available by provisioning additional cloud resources, ensuring scalability of the XR system to support numerous devices. In one example, many canonical maps can be maintained in memory for nearly instant access, or alternatively stored in high availability devices reducing system latency.

[0436] Further, performing localization for multiple devices in a cloud service may enable refinements to the process. Localization telemetry and statistics can provide information on which canonical maps to have in active memory and/or high availability storage. Statistics for multiple devices may be used, for example, to identify most frequently accessed canonical maps.

[0437] Additional accuracy may also be achieved as a result of processing in a cloud environment or other remote environment with substantial processing resources relative to a remote device. For example, localization can be made on higher density canonical maps in the cloud relative to processing performed on local devices. Maps may be stored in the cloud, for example, with more PCFs or a greater density of feature descriptors per PCF, increasing the accuracy of a match between a set of features from a device and a canonical map.

[0438] FIG. 61 is a schematic diagram of an XR system 6100. The user devices that display cross reality content during user sessions can come in a variety of forms. For example, a user device can be a wearable XR device (e.g., 6102) or a handheld mobile device (e.g., 6104). As discussed above, these devices can be configured with software, such as applications or other components, and/or hardwired to generate local position information (e.g., a tracking map) that can be used to render virtual content on their respective displays.

[0439] Virtual content positioning information may be specified with respect to global location information, which may be formatted as a canonical map containing one or more PCFs, for example. According to some embodiments, for example the embodiment shown in FIG. 61, the system 6100 is configured with cloud based services that support the functioning and display of the virtual content on the user device.

[0440] In one example, localization functions are provided as a cloud-based service 6106, which may be a micro-service. Cloud-based service 6106 may be implemented on any of multiple computing devices, from which computing resources may be allocated to one or more services executing in the cloud. Those computing devices may be interconnected with each other and accessible to devices, such as a wearable XR device 6102 and hand held device 6104. Such connections may be provided over one or more networks.

[0441] In some embodiments, the cloud-based service 6106 is configured to accept descriptor information from respective user devices and “localize” the device to a match-

ing canonical map or maps. For example, the cloud-based localization service matches descriptor information received to descriptor information for respective canonical map(s). The canonical maps may be created using techniques as described above that create canonical maps by merging maps provided by one or more devices that have image sensors or other sensors that acquire information about a physical world. However, it is not a requirement that the canonical maps be created by the devices that access them, as such maps may be created by a map developer, for example, who may publish the maps by making them available to localization service **6106**.

[0442] According to some embodiments, the cloud service handles canonical map identification, and may include operations to filter a repository of canonical maps to a set of potential matches. Filtering may be performed as illustrated in FIG. 29, or by using any subset of the filter criteria and other filter criteria instead of or in addition to the filter criteria shown in FIG. 29. In one embodiment, geographic data can be used to limit a search for matching canonical map to maps representing areas proximate to the device requesting localization. For example, area attributes such as Wi-Fi signal data, Wi-Fi fingerprint information, GPS data, and/or other device location information can be used as a coarse filter on stored canonical maps, and thereby limit analysis of descriptors to canonical maps known or likely to be in proximity to the user device. Similarly, location history of each device may be maintained by the cloud service such that canonical maps in the vicinity of the device's last location are preferentially searched. In some examples, filtering can include the functions discussed above with respect to FIGS. 31B, 32, 33, and 34.

[0443] FIG. 62 is an example process flow that can be executed by a device to use a cloud-based service to localize the device's position with canonical map(s) and receive transform information specifying one or more transformations between the device local coordinate system and the coordinate system of a canonical map. Various embodiments and examples may describe the one or more transforms as specifying transforms from a first coordinate frame to a second coordinate frame. Other embodiments include transforms from the second coordinate frame to the first coordinate frame. In yet other embodiments, the transforms enable transition from one coordinate frame to another, the resulting coordinate frames depend only on the desired coordinate frame output (including, for example, the coordinate frame in which to display content). In yet further embodiments, the coordinate system transforms may enable determination of a first coordinate frame from the second coordinate frame and the second coordinate frame from the first coordinate frame.

[0444] According to some embodiments, information reflecting a transform for each persistent pose defined with respect to the canonical map can be communicated to device.

[0445] According to one embodiment, process **6200** can begin at **6202** with a new session. Starting new session on the device may initiate capture of image information to build a tracking map for the device. Additionally, the device may send a message, registering with a server of a localization service, prompting the server to create a session for that device.

[0446] In some embodiments, starting a new session on a device optionally may include sending adjustment data from the device to the localization service. The localization ser-

vice returns to the device one or more transforms computed based on the set of features and associated poses. If the poses of the features are adjusted based on device-specific information before computation of the transformation and/or the transformations are adjusted based on device-specific information after computation of the transformation, rather than perform those computations on the device, the device specific information might be sent to the localization service, such that the localization service may apply the adjustments. As a specific example, sending device-specific adjustment information may include capturing calibration data for sensors and/or displays. The calibration data may be used, for example, to adjust the locations of feature points relative to a measured location. Alternatively or additionally, the calibration data may be used to adjust the locations at which the display is commanded to render virtual content so as to appear accurately positioned for that particular device. This calibration data may be derived, for example, from multiple images of the same scene taken with sensors on the device. The locations of features detected in those images may be expressed as a function of sensor location, such that multiple images yield a set of equations that may be solved for the sensor location. The computed sensor location may be compared to a nominal position, and the calibration data may be derived from any differences. In some embodiments, intrinsic information about the construction of the device may also enable calibration data to be computed for the display, in some embodiments.

[0447] In embodiments in which calibration data is generated for the sensors and/or display, the calibration data may be applied at any point in the measurement or display process. In some embodiments, the calibration data may be sent to the localization server, which may store the calibration data in a data structure established for each device that has registered with the localization server and is therefore in a session with the server. The localization server may apply the calibration data to any transformations computed as part of a localization process for the device supplying that calibration data. The computational burden of using the calibration data for greater accuracy of sensed and/or displayed information is thus borne by the calibration service, providing a further mechanism to reduce processing burden on the devices.

[0448] Once the new session is established, process **6200** may continue at **6204** with capture of new frames of the device's environment. Each frame can be processed to generate descriptors (including for example, DSF values discussed above) for the captured frame at **6206**. These values may be computed using some or all of the techniques described above, including techniques as discussed above with respect to FIGS. 14, 22 and 23. As discussed, the descriptors may be computed as a mapping of the feature points or, in some embodiments a mapping of a patch of an image around a feature point, to a descriptor. The descriptor may have a value that enables efficient matching between newly acquired frames/images and stored maps. Moreover, the number of features extracted from an image may be limited to a maximum number of features points per image, such as **200** feature points per image. The feature points may be selected to represent interest points, as described above. Accordingly, acts **6204** and **6206** may be performed as part of a device process of forming a tracking map or otherwise

periodically collecting images of the physical world around the device, or may be, but need not be, separately performed for localization.

[0449] Feature extraction at **6206** may include appending pose information to the extracted features at **6206**. The pose information may be a pose in the device's local coordinate system. In some embodiments, the pose may be relative to a reference point in the tracking map, such as a persistent pose, as discussed above. Alternatively or additionally, the pose may be relative to the origin of a tracking map of the device. Such an embodiment may enable the localization service as described herein to provide localization services for a wide range of devices, even if they do not utilize persistent poses. Regardless, pose information may be appended to each feature or each set of features, such that the localization service may use the pose information for computing a transformation that can be returned to the device upon matching the features to features in a stored map.

[0450] The process **6200** may continue to decision block **6207** where a decision is made whether to request localization. One or more criteria may be applied to determine whether to request localization. The criteria may include passage of time, such that a device may request localization after some threshold amount of time. For example, if localization has not been attempted within a threshold amount of time, the process may continue from decision block **6207** to act **6208** where localization is requested from the cloud. That threshold amount of time may be between ten and thirty seconds, such as twenty-five seconds, for example. Alternatively or additionally, localization may be triggered by motion of a device. A device executing the process **6200** may track its motion using an IMU and/or its tracking map, and initiate localization upon detection motion exceeding a threshold distance from the location where the device last requested localization. The threshold distance may be between one and ten meters, such as between three and five meters, for example. As yet a further alternative, localization may be triggered in response to an event, such as when a device creates a new persistent pose or the current persistent pose for the device changes, as described above.

[0451] In some embodiments, decision block **6207** may be implemented such that the thresholds for triggering localization may be established dynamically. For example, in environments in which features are largely uniform such that there may be a low confidence in matching a set of extracted features to features of a stored map, localization may be requested more frequently to increase the chances that at least one attempt at localization will succeed. In such a scenario, the thresholds applied at decision block **6207** may be decreased. Similarly, in an environment in which there are relatively few features, the thresholds applied at decision block **6207** may be decreased so as to increase the frequency of localization attempts.

[0452] Regardless of how the localization is triggered, when triggered, the process **6200** may proceed to act **6208** where the device sends a request to the localization service, including data used by the localization service to perform localization. In some embodiments, data from multiple image frames may be provided for a localization attempt. The localization service, for example, may not deem localization successful unless features in multiple image frames yield consistent localization results. In some embodiments, process **6200** may include saving feature descriptors and appended pose information into a buffer. The buffer may, for

example, be a circular buffer, storing sets of features extracted from the most recently captured frames. Accordingly, the localization request may be sent with a number of sets of features accumulated in the buffer. In some settings, a buffer size is implemented to accumulate a number of sets of data that will be more likely to yield successful localization. In some embodiments, a buffer size may be set to accumulate features from two, three, four, five, six, seven, eight, nine, or ten frames, for example). Optionally, the buffer size can have a baseline setting which can be increased responsive to localization failures. In some examples, increasing the buffer size and corresponding number of sets of features transmitted reduces the likelihood that subsequent localization functions fail to return a result.

[0453] Regardless of how the buffer size is set, the device may transfer the contents of the buffer to the localization service as part of a localization request. Other information may be transmitted in conjunction with the feature points and appended pose information. For example, in some embodiments, geographic information may be transmitted. The geographic information may include, for example, GPS coordinates or a wireless signature associated with the devices tracking map or current persistent pose.

[0454] In response to the request sent at **6208**, a cloud localization service may analyze the feature descriptors to localize the device into a canonical map or other persistent map maintained by the service. For example, the descriptors are matched to a set of features in a map to which the device is localized. The cloud based localization service may perform localization as described above with respect to device based localization (e.g., can rely on any of the functions discussed above for localization including, map ranking, map filtering, location estimation, filtered map selection, examples in FIGS. 44-46, and/or discussed with respect to a localization module, PCF and/or PP identification and matching etc.). However, instead of communicating identified canonical maps to a device (e.g., in device localization), the cloud-based localization service may proceed to generate transforms based on the relative orientation of feature sets sent from the device and the matching features of the canonical maps. The localization service may return these transforms to the device, which may be received at block **6210**.

[0455] In some embodiments, the canonical maps maintained by the localization service may employ PCFs, as described above. In such embodiments, the feature points of the canonical maps that match the feature points sent from the device may have positions specified with respect to one or more PCFs. Accordingly, the localization service may identify one or more canonical maps and may compute a transformation between the coordinate frame represented in the poses sent with the request for localization and the one or more PCFs. In some embodiments, identification of the one or more canonical maps is assisted by filtering potential maps based on geographic data for a respective device. For example, once filtered to a candidate set (e.g., by GPS coordinate, among other options) the candidate set of canonical maps can be analyzed in detail to determine matching feature points or PCFs as described above.

[0456] The data returned to the requesting device at act **6210** may be formatted as a table of persistent pose transforms. The table can be accompanied by one or more canonical map identifiers, indicating the canonical maps to which the device was localized by the localization service.

However, it should be appreciated that the localization information may be formatted in other ways, including as a list of transforms, with associated PCF and/or canonical map identifiers.

[0457] Regardless of how the transforms are formatted, at act 6212 the device may use these transforms to compute the location at which to render virtual content for which a location has been specified by an application or other component of the XR system relative to any of the PCFs. This information may alternatively or additionally be used on the device to perform any location based operation in which a location is specified based on the PCFs.

[0458] In some scenarios, the localization service may be unable to match features sent from a device to any stored canonical map or may not be able to match a sufficient number of the sets of features communicated with the request for the localization service to deem a successful localization occurred. In such a scenario, rather than returning transformations to the device as described above in connection with act 6210, the localization service may indicate to the device that localization failed. In such a scenario, the process 6200 may branch at decision block 6209 to act 6230, where the device may take one or more actions for failure processing. These actions may include increasing the size of the buffer holding feature sets sent for localization. For example, if the localization service does not deem a successful localization unless three sets of features match, the buffer size may be increased from five to six, increasing the chances that three of the transmitted sets of features can be matched to a canonical map maintained by the localization service.

[0459] Alternatively or additionally, failure processing may include adjusting an operating parameter of the device to trigger more frequent localization attempts. The threshold time between localization attempts and/or the threshold distance may be decreased, for example. As another example, the number of feature points in each set of features may be increased. A match between a set of features and features stored within a canonical map may be deemed to occur when a sufficient number of features in the set sent from the device match features of the map. Increasing the number of features sent may increase the chances of a match. As a specific example, the initial feature set size may be 50, which may be increased to 100, 150, and then 200, on each successive localization failure. Upon successful match, the set size may then be returned to its initial value.

[0460] Failure processing may also include obtaining localization information other than from the localization service. According to some embodiments, the user device can be configured to cache canonical maps. Cached maps permit devices to access and display content where the cloud is unavailable. For example, cached canonical maps permit device based localization in the event of communication failure or other unavailability.

[0461] According to various embodiments, FIG. 62 describes a high level flow for a device initiating cloud based localization. In other embodiments, various ones or more of the illustrated steps can be combined, omitted, or invoke other processes to accomplish localization and ultimately visualization of virtual content in a view of a respective device.

[0462] Further, it should be appreciated that, though the process 6200 shows the device determining whether to initiate localization at decision block 6207, the trigger for

initiating localization may come from outside the device, including from the localization service. The localization service, for example, may maintain information about each of the devices that is in a session with it. That information, for example, may include an identifier of a canonical map to which each device most recently localized. The localization service, or other components of the XR system, may update canonical maps, including using techniques as described above in connection with FIG. 26. When a canonical map is updated, the localization service may send a notification to each device that most recently localized to that map. That notification may serve as a trigger for the device to request localization and/or may include updated transformations, recomputed using the most recently sent sets of features from the device.

[0463] FIGS. 63A, B, and C are an example process flow showing operations and communication between a device and cloud services. Shown at blocks 6350, 6352 6354, and 6456 are example architecture and separation between components participating in the cloud based localization process. For example, the modules, components, and/or software that are configured to handle perception on the user device are shown at 6350 (e.g., 660, FIG. 6A). Device functionality for persisted world operations are shown at 6352 (including, for example, as described above and with respect to persisted world module (e.g., 662, FIG. 6A)). In other embodiments, the separation between 6350 and 6352 is not needed and the communication shown can be between processes executing on the device.

[0464] Similarly, shown at block 6354 is a cloud process configured to handle functionality associated with passable world/passable world modeling (e.g., 802, 812, FIG. 26). Shown at block 6356 is a cloud process configured to handle functionality associated with localizing a device, based on information sent from a device, to one or more maps of a repository of stored canonical maps.

[0465] In the illustrated embodiment, process 6300 begins at 6302 when a new session starts. At 6304 sensor calibration data is obtained. The calibration data obtained can be dependent on the device represented at 6350 (e.g., number of cameras, sensors, positioning devices, etc.). Once the sensor calibration is obtained for the device, the calibrations can be cached at 6306. If device operation resulted in a change in frequency parameters (e.g., collection frequency, sampling frequency, matching frequency, among other options) the frequency parameters are reset to baseline at 6308.

[0466] Once the new session functions are complete (e.g., calibration, steps 6302-6306) process 6300 can continue with capture of a new frame 6312. Features and their corresponding descriptors are extracted from the frame at 6314. In some examples, descriptors can comprise DSF's, as discussed above. According to some embodiments, the descriptors can have spatial information attached to them to facilitate subsequent processing (e.g., transform generation). Pose information (e.g., information, specified relative to the device's tracking map for locating the features in the physical world as discussed above) generated on the device can be appended to the extracted descriptors at 6316.

[0467] At 6318, the descriptor and pose information is added to a buffer. New frame capture and addition to the buffer shown in steps 6312-6318 is executed in a loop until a buffer size threshold is exceeded at 6319. Responsive to a determination that the buffer size has been met, a localiza-

tion request is communicated from the device to the cloud at **6320**. According to some embodiments, the request can be handled by a passable world service instantiated in the cloud (e.g. **6354**). In further embodiments, functional operations for identifying candidate canonical maps can be segregated from operations for actual matching (e.g., shown as blocks **6354** and **6356**). In one embodiment, a cloud service for map filtering and/or map ranking can be executed at **6354** and process the received localization request from **6320**. According to one embodiment, the map ranking operations are configured to determine a set of candidate maps at **6322** that are likely to include a device's location.

[0468] In one example, the map ranking function includes operations to identify candidate canonical maps based on geographic attributes or other location data (e.g., observed or inferred location information). For example, other location data can include Wi-Fi signatures or GPS information.

[0469] According to other embodiments, location data can be captured during a cross reality session with the device and user. Process **6300** can include additional operations to populate a location for a given device and/or session (not shown). For example, the location data may be stored as device area attribute values and the attribute values used to select candidate canonical maps proximate to the device's location.

[0470] Any one or more of the location options can be used to filter sets of canonical maps to those likely to represent an area including the location of a user device. In some embodiments, the canonical maps may cover relatively large regions of the physical world. The canonical maps may be segmented into areas such that selection of a map may entail selection of a map area. A map area, for example may be on the order of tens of meters squared. Thus, the filtered set of canonical maps may be a set of areas of the maps.

[0471] According to some embodiments, a localization snapshot can be built from the candidate canonical maps, posed features, and sensor calibration data. For example, an array of candidate canonical maps, posed features, and sensor calibration information can be sent with a request to determine specific matching canonical maps. Matching to a canonical map can be executed based on descriptors received from a device and stored PCF data associated with the canonical maps.

[0472] In some embodiments, a set of features from the device is compared to sets of features stored as part of the canonical map. The comparison may be based on the feature descriptors and/or pose. For example, a candidate set of features of a canonical map may be selected based on the number of features in the candidate set that have descriptors similar enough to the descriptors of the feature set from the device that they could be the same feature. The candidate set, for example, may be features derived from an image frame used in forming the canonical map.

[0473] In some embodiments, if the number of similar features exceeds a threshold, further processing may be performed on the candidate set of features. Further processing may determine the degree to which the set of posed features from the device can be aligned with the features of the candidate set. The set of features from the canonical map, like the features from the device, may be posed.

[0474] In some embodiments, features are formatted as a highly dimensional embedding (e.g., DSF, etc.) and may be compared using a nearest neighbor search. In one example,

the system is configured (e.g., by executing process **6200** and/or **6300**) to find the top two nearest neighbors using Euclidian distance, and may execute a ratio test. If the closest neighbor is much closer than the second closest neighbor, the system considers the closest neighbor to be a match. "Much closer" in this context may be determined, for example, by the ratio of Euclidean distance relative to the second nearest neighbor is more than a threshold times the Euclidean distance relative to the nearest neighbor. Once a feature from the device is considered to be a "match" to a feature in canonical map, the system may be configured to use the pose of the matching features to compute a relative transformation. The transformation developed from the pose information may be used to indicate the transformation required to localize the device to the canonical map.

[0475] The number of inliers may serve as an indication of the quality of the match. For example, in the case of DSF matching, the number of inliers reflects the number of features that were matched between received descriptor information and stored/canonical maps. In further embodiments, inliers may be determined in this embodiment by counting the number of features in each set that "match."

[0476] An indication of the quality of a match may alternatively or additionally be determined in other ways. In some embodiments, for example, when a transformation is computed to localize a map from a device, which may contain multiple features, to a canonical map, based on relative pose of matching features, statistics of the transformation computed for each of multiple matching features may serve as quality indication. A large variance, for example, may indicate a poor quality of match. Alternatively or additionally, the system may compute, for a determined transformation, a mean error between features with matching descriptors. The mean error may be computed for the transformation, reflecting the degree of positional mismatch. A mean squared error is a specific example of an error metric. Regardless of the specific error metric, if the error is below a threshold, the transformation may be determined to be usable for the features received from the device, and the computed transformation is used for localizing the device. Alternatively or additionally, the number of inliers may also be used in determining whether there is a map that matches a device's positional information and/or descriptors received from a device.

[0477] As noted above, in some embodiments, a device may send multiple sets of features for localization. Localization may be deemed successful when at least a threshold number of sets of features match, with an error below a threshold, and/or a number of inliers above a threshold, a set of features from the canonical map. That threshold number, for example, may be three sets of features. However, it should be appreciated that the threshold used for determining whether a sufficient number of sets of feature have suitable values may be determined empirically or in other suitable ways. Likewise, other thresholds or parameters of the matching process, such as degree of similarity between feature descriptors to be deemed matching, the number of inliers for selection of a candidate set of features, and/or the magnitude of the mismatch error, may similarly be determined empirically or in other suitable ways.

[0478] Once a match is determined, a set of persistent map features associated with the matched canonical map or maps is identified. In embodiments in which the matching is based on areas of maps, the persistent map features may be the map

features in the matching areas. The persistent map features may be persistent poses or PCFs as described above. In the example of FIG. 63, the persistent map features are persistent poses.

[0479] Regardless of the format of the persistent map features, each persistent map feature may have a predetermined orientation relative to the canonical map in which it is a part. This relative orientation may be applied to the transformation computed to align the set of features from the device with the set of features from the canonical map to determine a transformation between the set of features from the device and the persistent map feature. Any adjustments, such as might be derived from calibration data, may then be applied to this computed transformation. The resulting transformation may be the transformation between the local coordinate frame of the device and the persistent map feature. This computation may be performed for each persistent map feature of a matching map area, and the results may be stored in a table, denoted as the `persistent_pose_table` in 6326.

[0480] In one example, block 6326 returns a table of persistent pose transforms, canonical map identifiers, and number of inliers. According to some embodiments, the canonical map ID is an identifier for uniquely identifying a canonical map and a version of the canonical map (or area of a map, in embodiments in which localization is based on map areas).

[0481] In various embodiments, the computed localization data can be used to populate localization statistics and telemetry maintained by the localization service at 6328. This information may be stored for each device, and may be updated for each localization attempt, and may be cleared when the device's session ends. For example, which maps were matched by a device can be used to refine map ranking operations. For example, maps covering the same area to which the device previously matched may be prioritized in the ranking. Likewise, maps covering adjacent areas may be given higher priority over more remote areas. Further, the adjacent maps might be prioritized based on a detected trajectory of the device over time, with map areas in the direction of motion being given higher priority over other map areas. The localization service may use this information, for example, upon a subsequent localization request from the device to limit the maps or map areas searched for candidate sets of features in the stored canonical maps. If a match, with low error metrics and/or a large number or percentage of inliers, is identified in this limited area, processing of maps outside the area may be avoided.

[0482] Process 6300 can continue with communication of information from the cloud (e.g., 6354) to the user device (e.g., 6352). According to one embodiment, a persistent pose table and canonical map identifiers are communicated to the user device at 6330. In one example, the persistent pose table can be constructed of elements including at least a string identifying a persistent pose ID and a transform linking the device's tracking map and the persistent pose. In embodiments in which the persistent map features are PCFs the table may, instead, indicate transformations to the PCFs of the matching maps.

[0483] If localization fails at 6336, process 6300 continues by adjusting parameters that may increase the amount of data sent from a device to the localization service to increases the chances that localization will succeed. Failure, for example, may be indicated when no sets of features in the

canonical map can be found with more than a threshold number of similar descriptors or when the error metric associated with all transformed sets of candidate features is above a threshold. As an example of a parameter that may be adjusted, the size constraint for the descriptor buffer may be increased (of 6319). For example, where the descriptor buffer size is five, localization failure can trigger an increase to at least six sets of features, extracted from at least six image frames. In some embodiments, process 6300 can include a descriptor buffer increment value. In one example, the increment value can be used to control the rate of increase in the buffer size, for example, responsive to localization failures. Other parameters, such as parameters controlling the rate of localization requests, may be changed upon a failure to find matching canonical maps.

[0484] In some embodiments, execution of 6300 can generate an error condition at 6340, which includes execution where the localization request fails to work, rather than return a no match result. An error, for example, may occur as a result of a network error making the storage holding a database of canonical maps unavailable to a server executing the localization service or a received request for localization services containing incorrectly formatted information. In the event of an error condition, in this example, the process 6300 schedules a retry of the request at 6342.

[0485] When a localization request is successful, any parameters adjusted in response to a failure may be reset. At 6332, process 6300 can continue with an operation to reset frequency parameters to any default or baseline. In some embodiments 6332 is executed regardless of any changes thus ensuring baseline frequency is always established.

[0486] The received information can be used by the device at 6334 to update a cache localization snapshot. According to various embodiments, the respective transforms, canonical maps identifiers, and other localization data can be stored by the device and used to relate locations specified with respect to the canonical maps, or persistent map features of them such as persistent poses or PCFs to locations determined by the device with respect to its local coordinate frame, such as might be determined from its tracking map.

[0487] Various embodiments of processes for localization in the cloud can implement any one or more of the preceding steps and be based on the preceding architecture. Other embodiments may combine various ones or more of the preceding steps, execute steps simultaneously, in parallel, or in another order.

[0488] According to some embodiments, localization services in the cloud in the context of cross reality experiences can include additional functionality. For example, canonical map caching may be executed to resolve issues with connectivity. In some embodiments, the device may periodically download and cache canonical maps to which it has localized. If the localization services in the cloud are unavailable, the device may run localizations itself (e.g., as discussed above—including with respect to FIG. 26). In other embodiments, the transformations returned from localization requests can be chained together and applied in subsequent sessions. For example, a device may cache a train of transformations and use the sequence of transformations to establish localization.

[0489] Various embodiments of the system can use the results of localization operations to update transformation information. For example, the localization service and/or a device can be configured to maintain state information on a

tracking map to canonical map transformations. The received transformations can be averaged over time. According to one embodiment, the averaging operations can be limited to occur after a threshold number of localizations are successful (e.g., three, four, five, or more times). In further embodiments, other state information can be tracked in the cloud, for example, by a passable world module. In one example, state information can include a device identifier, tracking map ID, canonical map reference (e.g., version and ID), and the canonical map to tracking map transform. In some examples, the state information can be used by the system to continuously update and get more accurate canonical map to tracking map transforms with every execution of the cloud-based localization functions.

[0490] Additional enhancements to cloud-based localization can include communicating to devices outliers in the sets of features that did not match features in the canonical maps. The device may use this information, for example, to improve its tracking map, such as by removing the outliers from the sets of features used to build its tracking map. Alternatively or additionally, the information from the localization service may enable the device to limit bundle adjustments for its tracking map to computing adjustments based on inlier features or to otherwise impose constraints on the bundle adjustment process. According to another embodiment, various sub-processes or additional operations can be used in conjunction and/or as alternatives to the processes and/or steps discussed for cloud based localization. For example, candidate map identification may include accessing canonical maps based on area identifiers and/or area attributes stored with respective maps.

Deep Correspondences

[0491] Described herein are methods and apparatus for efficiently and accurately finding matching sets of feature points, such as may occur in localizing XR devices in large scale environments in real-time. Accordingly, matching sets of features as part of localization is described herein to illustrate techniques that may lead to fast and accurate matches. Some or all of these techniques may be applied when matching sets of features in other contexts, such as when searching for a match between a portion of a tracking map and a canonical map as part of a map merge process.

[0492] Localizing XR devices may require making comparisons to find a match between a set of 2D features from one or more images captured by the XR devices and a set of feature points, which may be 3D map points in a stored canonical map. Maps for large scale environments may include a large number of 3D map points.

[0493] Some of the 3D map points may be captured at different times during the day or at different seasons compared with the 2D image features. The different dimensionality, different lighting condition and other conditions makes it more difficult to accurately find matching sets of features. Accurate localization, for example, in large and very large scale environments, may require a larger number of sets of 2D features to be compared to provide an accurate localization result. As a result, localizing XR devices in large and very large scale environments takes more time and consumes more computing power, causing delays in displaying virtual contents and affecting the realism of XR experiences.

[0494] The inventors have recognized and appreciated methods and apparatus that localize XR devices in large and

very large scale environments with reduced time and improved accuracy with a service that searches for matching sets of features using subsets of features with matching descriptors. The system may include a component that assess the likelihood that including a pair of features with matching descriptions in a subset will lead to finding matching sets of features.

[0495] In some embodiments, a localization service guided in the selection of subsets of features with matching features by the component may provide a localization result in real time, such as no more than ten milliseconds, five milliseconds, or two milliseconds in some embodiments. In some embodiments, the localization service guided by the component may reduce the runtime to respond to a localization request by about ten times, for example, from 25 ms to 3 ms with one hundred corresponding features, with similar or improved localization accuracy. In some embodiments, the localization service guided by the component may reduce the number of iterations run by the pose estimation algorithms, which may determine a transformation that aligns one subset of features with another subset of features with matching descriptors, by about ten times, for example, from one hundred iterations to twelve iterations, with similar or improved localization accuracy.

[0496] The localization service may be on the XR devices, on the cloud, or both. In some embodiments, a persisted map may be downloaded to an XR device for localizing the XR devices in the map, for example, as described above with respect to FIG. 26. In some embodiments, an XR device may upload information (e.g., 2D features and/or associated metadata) to a cloud containing the localization service and receive its location in one or more persisted maps, for example, as described above with respect to FIGS. 61-63C.

[0497] Such a localization service may receive and apply any one or more types of location metadata in connection with a localization request to select a set of candidate maps, frame descriptors or other criteria. These criteria may be used to select one or more canonical maps or segments of a canonical map against which localization may be attempted. Other criteria, such as deep key frame descriptors, may be used to further down select from the set of candidate maps or to identify segments of the candidate maps against which to attempt localization.

[0498] Following such down selection, a feature-level comparison may be performed. A set of features in a key frame generated by the device may be compared to a set of features in an identified segment of a candidate map, for example. Where there is more than one candidate map or more than one identified segment of a candidate map, the feature set from the device may be compared to multiple sets of features before a localization result is determined. The comparisons may continue until all candidate segments of all candidate maps are processed. Localization may be deemed successful if a set of features from a candidate map that best corresponds to the feature set from the device has an error below some threshold. Alternatively or additionally, processing may complete once a set of features from a candidate map with a correspondence with an error below some threshold is identified. Regardless, the localization process may involve comparing multiple sets of features. For simplicity, processing of one set is described.

[0499] The localization service may respond to requests from one or more XR devices to localize with respect to a set of one or more persisted maps. The request may include

2D features extracted from images of the physical world around the device. In some embodiments, the images may be captured by multiple sensors, such as cameras, of the device. In some embodiments, a set of features from an XR device may be based on information captured by multiple sensors simultaneously, which may increase the speed and/or accuracy of finding matching sets of features.

[0500] Those 2D features may be posed relative to coordinate frames used by the XR devices in a way that the 2D features captured by multiple sensors can be processed together to provide the localization result. In some embodiments, the localization service may integrate 2D features from multiple sensors by embedding sensor extrinsic parameters in the 2D features. The sensor extrinsic parameters may include the physical dimensions of the sensors, the distances between the sensors, the physical dimensions of an XR device display, the locations of the sensors on the XR device, etc. Appropriate transformations to represent the sensor extrinsic parameters, which may be used to implement the embedding, may be determined through a calibration process.

[0501] This set of 2D features may serve as an input to the localization process. The localization service may identify a set of 3D features from a candidate map. The set of 3D features may be selected based on feature descriptors that match descriptors of the 2D features. The inventors have recognized and appreciated that some of the matched correspondences may be true, meaning that the corresponding features in the 2D and 3D sets of features represent the same features in the physical world, while others of the matched correspondences may be false, meaning that the corresponding features in the 2D and 3D sets of features, despite having matching descriptors, do not represent the same features in the physical world. Accordingly, assessing the quality of the correspondence, and weighting the process towards consideration of corresponding features of high quality may shorten computing time and/or increase the accuracy of the localization result.

[0502] The localization service may include a component configured to assess the quality of the corresponding features. The localization service may then identify a transformation that aligns the 2D set of features to the 3D sets of features, using processing on subsets of corresponding features, with inclusion in the subsets weighted towards using pairs of corresponding features with high quality.

[0503] In some embodiments, searching for a transformation between two sets of features may be computationally intensive. In some embodiments, the localization service may use gravity orientation to reduce the degrees of freedom that should be searched when matching 2D features to 3D features. Both the 2D features from the device and the 3D features from a stored map may be expressed in a coordinate frame with one coordinate aligned with gravity. In some embodiments, the localization service may, prior to searching for a transformation, rotate the set of 2D features such that the coordinate of the 2D feature set aligned with gravity aligns with the coordinate of the set of 3D features that is also aligned with gravity. Alternatively or additionally, any searching for transformations may exclude transformations that would change the alignment of the feature set relative to gravity.

[0504] From this transformation, the localization service may compute and return to the device a transformation to relate its local coordinate frame in which the 2D features

were posed to a coordinate frame of the persisted map from which the set of 3D features was obtained.

[0505] The component configured to assess the quality of the matched correspondences may include an artificial neural network. The neural network may be trained to provide a quality metric for each pair of corresponding features. The quality metric may indicate the likelihood that a 2D-3D feature pair identify the same feature in the physical world. In some embodiments, the quality metric may be a probability in the range of zero to one, for example, with one indicating a correct match and zero indicating a false match.

[0506] The neural network may be trained with a data set that has feature points represented as both 2D and 3D features. The training set may also include 2D and 3D features that represent different features in the physical world but have the same or similar descriptors. Such a data set may be generated synthetically, such as from computer graphics generated data depicting environments in which XR devices may operate. The data may have noise applied to it, indicative of noise that may exist in the measurement process or be otherwise distorted to be representative of real-world data.

[0507] Training of the neural network may be based on loss function that penalizes wrong results. A wrong result may be either assigning a low quality to a pair of 2D and 3D features that represent the same feature in the physical world or assigning a high quality to a pair of 2D and 3D features that represent different features in the physical world. Alternatively or additionally, the loss function may promote correct results, which may be either assigning a high quality to a pair of 2D and 3D features that represent the same feature in the physical world or assigning a low quality to a pair of 2D and 3D features that represent different features in the physical world.

[0508] FIG. 64 is a block diagram illustrating a portion of an XR system 6400 that provides large and very large scale localization, according to some embodiments. FIG. 65 is a schematic diagram illustrating information about a physical world being processed by the XR system 6400, according to some embodiments. Processing as shown in FIG. 64 may be implemented on a portable device or may be distributed across a device and one or more remote processing systems, such as a cloud service.

[0509] The XR system 6400 may include one or more sensors 6402 configured to capture information 6404 about a physical world. In some embodiments, the sensors may include one or more image sensors, for example, cameras 552 and/or 553 in FIG. 5B, which may output grayscale and/or color image frames at fixed time intervals. In some embodiments, the sensors may include one or more inertial measurement units (IMU), for example, IMU 557 in FIG. 5B, which may detect movements and orientations. The image frames output by the cameras may be appended with orientations captured by the IMU such as gravity orientations of the device when the images are taken. The images may be rotated such that the gravity orientations are aligned, which reduces the degrees of freedom to search for a localization service.

[0510] FIG. 65 illustrates embedding of data collected by the sensors 6402. In the example of FIG. 65, four cameras cam0-cam3 of an XR device are illustrated. The XR device may have a device coordinate frame 6502. In some embodiments, the device coordinate frame 6502 may be a keyrig coordinate frame, indicating a pose of the device when a

keyrig (e.g., keyrig 704 FIG. 7) is captured. In some embodiments, the pose of the keyrig may be indicated in a coordinate frame of a tracking map built on the XR device, for example, the world coordinate frame 86 (FIGS. 9, 10). In some embodiments, the pose of the keyrig may be determined based on the location of a display of the device such that virtual content can be displayed correctly on the display once the device is localized to a map.

[0511] The XR device may attempt to localize to one or more maps persisted in a database 6410, for example, a canonical map 120 in FIG. 16. The persisted maps may include map points representing 3D features of the physical world, for example, map points 3802 in FIG. 38A. The map points may share a map coordinate frame and be referable through the shared map coordinate frame, for example, a canonical coordinate frame 4806C in FIG. 39A.

[0512] The system 6400 may provide a pose estimation 6424, which may be in the form of a transformation between a device coordinate frame and a map coordinate frame. In the example of FIG. 65, an example of the pose estimation 6424 is illustrated as a transformation 6506 between the device coordinate frame 6502 and the map coordinate frame 6504. The transformation 6506 may include translations and rotations between the coordinate frames 6502 and 6504, which may be represented by a matrix.

[0513] As described above, a pose may be estimated by finding a transformation that aligns the two sets of features. The XR system 6400 may include a feature extraction component 6406 configured to identify features from sensor-captured information and output feature information 6408. As described above, examples of features may include corners and edges of objects in a physical world. In the example of FIG. 65, the feature extraction component 6406 may receive an image 6508 captured by cam3, and identify one or more features 6510k in the image 5408. The feature extraction component 6406 may also receive images from cam0, cam1, and cam2, and identify features such as 6510i and 6510j. Descriptors for each feature may also be generated to enable efficient matching. In some embodiments, the feature information 6408 may include the descriptors (e.g., DSF descriptors in FIG. 25).

[0514] The feature extraction component 6406 may also be configured to append sensor extrinsic parameters to the identified features such that the XR system 6400 can simultaneously process information captured by multiple sensors. In some embodiments, the feature information 6408 may include a 6d vector by combining two 3d vectors. A first 3d vector may indicate the position of the feature in a sensor coordinate frame of the sensor that captured the image containing the feature (e.g., $u_i v_i w_i, u_j v_j w_j, u_k v_k w_k$). A second 3d vector may indicate the position in the device coordinate frame 6502 of the sensor that captured the image containing the feature (e.g., t_0, t_1, t_2, t_3). For example, the feature information 6408 for the feature 6510k may include a unique descriptor for the feature, the first vector $u_k v_k w_k$, and the second vector t_3 .

[0515] In some embodiments, the feature extraction component 6406 may be on the devices. The feature information 6408 may be communicated to a localization service 6426. Communicating the feature information 6408 may consume less bandwidth than sending full images. However, the present application is not limited in this regard. In some embodiments, the localization service 6426 may include the feature extraction component 6406.

[0516] The localization service 6426 may include a feature matching component 6414, which may receive feature information 6408 from the feature extraction component 6406. The feature matching component 6414 may also receive map information 6412 from one or more maps persisted in the database 6410. The one or more maps may be selected from the database 6410 by the method of FIG. 28 and/or the method 900 of FIG. 29. The map information 6412 may include map points representing 3D features of the physical world. The feature matching component 6414 may be configured to match corresponding features from feature information 6408 and map information 6412 by, for example, selecting pairs of features from the 2D and 3D sets that have the same or similar descriptors.

[0517] In the example of FIG. 65, the feature matching component 6414 may determine that features 6510i, 6510j, 6510k match with map points 6512i, 6512j, 6512k, respectively. The map points 6512i, 6512j, 6512k may be expressed as 3d vectors $x_i y_i z_i, x_j y_j z_j, x_k y_k z_k$, indicating the positions of the map points in the map coordinate frame 6504. As illustrated, some matched correspondences may be correct, for example, the correspondence between 6510j and 6512j, and the correspondence between 6510k and 6512k; some matched correspondences may be false, for example, the correspondence between 6510i and 6512i.

[0518] The localization service 6426 may include a matched correspondences quality prediction component 6418 configured to assess the quality of the pairs of matched features 6416 provided by the feature matching component 6414. The matched correspondences quality prediction component 6418 may provide quality information 6420 to a pose estimation component of the localization service 6426 such that the pose estimation component 6422 can be guided by the quality information 6420 when computing the pose estimation 6424.

[0519] The quality information 6420 may include quality metrics for each pair of matched features. In some embodiments, the quality metric may be a probability in the range of zero to one, with a higher assigned quality indicating a higher likelihood that the matched features represent the same location in the physical world. It should be appreciated that a component may be configured to indicate the quality of the matched correspondences with other numeric values, for example, indicating a likely correct match with zero and likely false match with one, or vice versa.

[0520] The pose estimation component 6422 may process the pairs of matched features 6416, output by the feature matching component 6414, based on the quality information 6420, output by the matched feature quality prediction component 6418. In some embodiments, the pose estimation component 6422 may select a first subset of pairs of matched features from the set 6416, and compute a first pose based on the selected subset. The number of correspondences in the subset may be five, seven, eight, or any value that is sufficient to provide a valid transformation matrix. The selection of the first subset may be guided by the quality information 6420, with the selection weighted towards higher quality matching features. For example, the pose estimation component 6422 may randomly select the first subset of matching features from all pairs of matching features, but with a biased selection such that matching features having higher quality metrics are more likely to be

selected than those with low quality metrics. The likelihood of a pair of matched features being selected may be proportional to the quality metric.

[0521] The pose estimation component **6422** may compute a transformation that aligns the 2D features of the selected subset of pairs of matched features with corresponding 3D features in the subset. In systems that consider gravity, transformations that alter the orientation of the 2D features relative to gravity may not be considered.

[0522] The pose estimation component **6422** may determine the accuracy of the first pose by applying the computed transformation to a larger set of pairs of matched features, including some or all of the pairs of matched features not included in the first subset. In some embodiments, the pose estimation component **6422** may determine accuracy by computing projection errors for individual correspondences. A projection error for a matched correspondence may indicate a distance between the position of the 2D feature in the image containing the 2D feature and the position of the matched map point projected to a plane that the image extends.

[0523] In some scenarios, pose estimation component **6422** may compute further estimated poses to ensure an accurate pose is computed. In some embodiments, the pose estimation component **6422** may determine the estimated pose is accurate when the number of pairs of matched features having projection errors below a threshold error (e.g., two pixels) is above a threshold value (e.g., 80% of the set **6416**). This check may be made as the first pose estimation is completed. Alternatively, the pose estimation component **6422** may compute a plurality of pose estimations and check the accuracy of the best of the pose estimations relative to a threshold value.

[0524] When the pose estimation component **6422** determines that the first localization is not sufficiently accurate, or in embodiments in which multiple pose estimations are computed before assessing the suitability of the best one, the pose estimation component **6422** may randomly select a second subset of pairs of matched features from the set **6416**. The selection of a further subset of pairs of matched features, as for the first subset, may be weighted by the quality information **6420**. The pose estimation component **6422** may compute a second poses to align the 2D features of the selected second subset and the 3D features. This process may be continued, with additional subsets being selected and processed, until an accurate localization result is found, and/or other stop conditions are reached such as processing of a predetermined number of subsets. It should be appreciated that with the guidance from the quality information **6420**, a pose estimation component may be able to provide an accurate localization result with selecting and computing based on ten times fewer number of subsets of matched correspondences than without the guidance.

[0525] In some embodiments, the matched feature quality predication component **6418** may include an artificial neural network **6602**, for example, as illustrated in a subsystem **6600** of the XR system **6400** in FIG. 66, according to some embodiments. The neural network **6602** may include an encoding layer **6604**, a decoding layer **6616**, and an intermediate layer **6622** between the encoding layer **6604** and the decoding layer **6616**.

[0526] The encoding layer **6604** may include multi-layer perceptrons (MLP) **6606**. Each MLP **6606** in the encoding layer **6604** may receive one or more pairs of matched

features from the set **6416**. In some embodiments, a matched pair of features may include a 9d vector, which may be a combination of the 6d vector of a 2D feature as described above (e.g., $u_k v_k w_k$ & t_3) and the 3d vector indicating the position of the matched 3D feature in the map (e.g., $x_k y_k z_k$). The MLPs may provide encoded vectors **6608**, which may provide a higher dimensionality of features than the input correspondence. In some embodiments, the encoded vectors **6608** may include a 64d vector.

[0527] The intermediate layer **6622** may include one or more residual network blocks **6610** connected by element-wise summation blocks **6614**. A residual network block **6610** may include MLPs, and a sub-block configured to normalize a distribution of outputs of the MLPs of the residual network block.

[0528] The decoding layer **6616** may also include MLPs. Each MLP in the decoding layer **6616** may receive encoded output from the intermediate layer **6622**, and output a decoded vector **6618**, which may have a same dimension as the input correspondence. The decoded vectors **6618** may be converted into the quality information **6420** through activation blocks **6620**.

[0529] FIG. 67 is a flow chart illustrating a method **6700** of generating a reference dataset for training the neural network **6602**, according to some embodiments. The method **6700** may start by creating (Act **6702**) a reference dataset comprising 2D-3D feature correspondences. The reference dataset may be synthetic or real-world data. In some embodiments, the reference data set may include pairs of features that correspond to correct matches and those that correspond to incorrect matches.

[0530] At Act **6704**, the method **6700** may compute a pose based on the reference dataset by, for example, using the pose estimation component **6420**. At Act **6706**, the estimated pose may be used to compute ground truth (GT) projection errors for individual correspondences. In scenarios in which synthetic data is used, ground truth may be determined from any transformation between the 2D and 3D sets introduced in forming the synthetic data. For other data sets, ground truth may be determined in other ways, including as a result of manual review of data sets or the images from which the data sets were generated. For a 2D-3D feature correspondence, a GT projection error may indicate an actual distance between the position of the 2D feature of the correspondence on an image containing the 2D feature and the position of a 2D feature that corresponds to the 3D feature in the correspondence.

[0531] At Act **6708**, the method **6700** may compute GT weights for the correspondences of the reference dataset based on the computed GT projection errors. In some embodiments, the GT weights W_{GT} may be defined as the Cauchy weight function of the GT projection error as shown below:

$$W_{GT} = \frac{C^2}{C^2 + \text{error}^2}$$

where C may have the value of the error for a GT weight of 0.5. In some embodiments, the training dataset **6710** may include the 2D-3D feature correspondences of the reference dataset and the computed GT weights.

[0532] FIG. 68 is a flow chart illustrating a method **6800** of training the neural network **6602**, according to some

embodiments. The method **6800** may start by computing (Act **6802**) quality metrics for each pair of features in the training dataset using the neural network **6602**. At Act **6804**, the method **6800** may compute a loss based on the difference between the quality metrics and the GT weights. In some embodiments, the loss may be a regression loss on GT weights W_{GT} , for example, a mean square error between the quality metrics and corresponding GT weights for the pair of features in the training set **6710**. At Act **6806**, the method **6800** may modify the weights of the neural network **6602** based on the loss so as to reduce the loss. Pairs of matching features in the training set may be applied successively, updating the neural network as each pair is processed, so as to decrease the loss.

[0533] Regardless of the specific method by which the neural network is trained, once trained, it may be used to guide the selection of subsets of the matching pairs of features to reduce the number of subsets processed to identify a pose between two sets of features, as described above in connection with FIG. **64**. This process may be preceded or followed by other processing. One or more criteria may be applied before the processing illustrated in FIG. **64** to determine which sets of 3D features to compare to a set of 2D features. If multiple sets of 3D features are to be compared, the process of FIG. **64** may be repeated for each set of 3D features to determine the best matching feature set. Following the processing of FIG. **64**, identified pose may be converted to a localization result.

[0534] The process for pose estimation of FIG. **64** alternatively or additionally may be applied in other contexts. For example in merging a tracking map to a canonical map, a set of 2D features from the tracking map, such as those associated with a persistent pose, may be compared to 3D features in the canonical map. The pose estimation process of FIG. **64** may be applied to determine whether sets of features match and, if so, a transformation between the sets of features. Computation of that transformation may be used to align the tracking map to the canonical map so that the maps may be combined.

[0535] Accordingly, it should be appreciated that the process of pose estimation as described herein in connection with localizing a device may be applied in other scenarios.

Gravity Preserving Map Merging

[0536] As described above, a set of canonical maps of the environment in which multiple XR devices operate may provide an enhanced XR experience to multiple users using different computing devices (such as, portable devices including smart phones and XR devices) by enabling their local representations of the world (e.g., in the form of tracking maps) to be transformed to the same coordinate frame. By localizing to the same frame of reference, all of the devices may present to their users virtual content in the same location relative to the physical world, creating a more realistic and immersive shared user experience.

[0537] The canonical maps may be created by incrementally merging tracking maps received from the users' devices to create a virtual representation of the world. In some embodiments, users' devices may send their tracking maps to a cloud-based service to merge with map(s) of the environment previously stored in the cloud-based service. The map merge component of the cloud service may attempt to identify a portion of each new tracking map that represents the same region of the physical world as a portion of

an existing canonical map. This identification may be made in steps, with an initial step involving finding likely regions of overlap by matching location metadata associated with regions of the already stored canonical maps to similar location metadata associated with the tracking map. Alternatively or additionally, descriptors computed for key frames in the tracking map may be matched to similar descriptors associated with portions of the canonical map. These likely regions of overlap serve as candidates for further processing.

[0538] These candidate regions of overlap may be further processed to identify feature points in the tracking map with corresponding feature points in the canonical map. Corresponding features may be identified based on characteristics of the features. For example, corresponding features may be identified based on similar descriptors, such as DSF descriptors described above, being assigned to features in the tracking map and the canonical map.

[0539] Further steps in the map merge process may entail determining a transformation between the sets of corresponding features that provides a suitably low error between the positions of the transformed features in a first set of features derived from the tracking map and a second set of features derived from the canonical map. One or more error metrics may be applied to compare error to a threshold. For example, the number or percentage of feature points in the tracking map that, after transformation, are within a threshold distance of their corresponding feature point in the canonical map. Alternatively or additionally, error may be measured based on distance between transformed feature points from the tracking map and their corresponding features in the canonical map. As a specific example, an error measure may be the root mean squared distance for the set of corresponding feature points.

[0540] Multiple transformations may be attempted until a transformation with suitable error is identified. Guided search strategies may be used to select the transformations tested during a search for a transformation that provides a low error meeting an ending criteria for the search. For example, searching for a suitable transformation may entail performing a pose estimation on weighted subsets of pairs of matching features and then checking the error when the estimated pose is applied as a transformation to the points of the tracking map. Such a search may use techniques as described above in connection with FIG. **64**.

[0541] Regardless of the specific search strategy, if the search process is concluded with no suitable transformation identified, the merge process may continue with selecting another candidate region of one of the canonical maps to compare to the same region of the tracking map. Alternatively or additionally, a different region of the tracking map may be selected for comparison to candidate regions of the canonical maps in an attempt to find a transformation with suitably low error.

[0542] If all candidate regions are searched without identifying a suitable transformation, the tracking map may not be combined with any of the existing maps. One or more other actions might be taken, such as ending the merge process with no action. In some embodiments, the tracking map may be deemed to represent a portion of the physical world for which no canonical map already exists, and the tracking map may be stored as a new canonical map in the set of canonical maps. In this way, an initial canonical map

inherits properties from the tracking map, including its orientation with respect to gravity, as described above.

[0543] Conversely, the ability to determine a transformation with suitably low error serves as confirmation that the candidate region likely represents the same region of the physical world as the tracking map and the two maps can be merged. In some embodiments, identification of a suitable transformation for merging two maps may entail additional processing, such as finding multiple portions of the tracking map that can be aligned, with sufficiently close transformations. For example, a tracking map may be merged with a canonical map upon finding multiple key frames in the tracking map that have similar poses with respect to the canonical map. In that scenario, the determined transformation may be a combination of the transformations determined for the matching regions, such as an average of the transformations.

[0544] Regardless of the specific approach by which a suitable transformation is determined, that transformation may be applied to the entire tracking map, which may align the tracking map to the canonical map. Overlapping pieces may be combined, such as by averaging or aggregating feature points in the two maps or overwriting feature points of the canonical maps with feature points from the new map. Non-overlapping portions of the tracking map may be appended to the canonical map.

[0545] The inventors have recognized that errors may be introduced during the map merging process when there are errors in the determined transformation between a tracking map received from the user's device and a canonical map that the tracking map is combined with. As a result of errors in the transformations, the tracking maps as merged into the canonical maps may be skewed, causing the map merging process to output a skewed merged map. Such alignment failures may be caused for diverse reasons, such as, confusing visual similarities between feature points. The inventors have developed a merging technique (performed by map merge portion 810 of FIG. 26, for example) that reduces merging errors by ensuring that, after applying a determined transformation, a gravity direction of the tracking map aligns with a gravity direction of the canonical map with which the tracking map is to be merged.

[0546] In some embodiments, the gravity direction of a tracking map is obtained or derived from an inertial measurement unit (e.g., IMU 557) included in the portable device. As a specific example, the gravity direction of the tracking map may be obtained or derived from one or more accelerometers (e.g., 3D accelerometers) included in one or more IMUs. The tracking map may be generated at the portable device based on sensor data collected from various sensors on the portable device, including for example cameras that are mechanically coupled to the same support structure as the one or more IMUs. As images are output from the cameras, they may be transformed so as to have a predetermined orientation with respect to gravity. As features from these images are combined into a tracking map, those features will have the same orientation with respect to gravity, such that all generated tracking maps may have the same alignment with respect to gravity.

[0547] As all tracking maps may have the same alignment with respect to gravity, any tracking map sent to the cloud for map merge processing will be aligned with respect to gravity. When a tracking map does not match any existing canonical map and is used as the start of a new canonical

map, that canonical map may inherit properties from the tracking map, including its orientation with respect to gravity. Accordingly, the canonical maps may each start aligned with gravity. In systems in which canonical maps are initiated other than from a tracking map, the canonical map formation process may be restricted otherwise to start each canonical map aligned with gravity.

[0548] In some scenarios, as a result of an error in a determined transformation, the tracking map as transformed by the cloud-based service to align a set of features in the tracking map with matching features in the canonical map may be skewed such that it no longer aligns with gravity. Despite a low value of an error metric derived from closeness of match between features, such a condition signals an error has occurred in the alignment process.

[0549] By enforcing a gravity constraint such that transformed tracking maps that are skewed with respect to gravity are not used for merging, the accuracy of merged maps is improved. Such a constraint may be implemented in the map merging process by comparing the gravity direction of the transformed tracking map with a gravity direction of an environment map in the cloud, and merging the two maps only when the gravity directions align. This constraint may be applied at one or more stages in the map merge process. For example, once a transformation is determined from aligning matching feature points, the effect of that transformation on the orientation of the tracking map with respect to gravity may be assessed. A transformation that rotates the tracking map with respect to gravity by more than a threshold amount may be discarded, and the process of searching for a suitable transformation to align the tracking map with a portion of a canonical map may be resumed or may end without merging the tracking map. Alternatively or additionally, as part of a search for a suitable transformation in which candidate transformations are tested on subsets of features, only candidate transformations that do not rotate the subset of features with respect to gravity may be considered.

[0550] FIG. 69 illustrates a map merge process that ensures that a gravity direction of a tracking map aligns with a gravity direction of an environment map with which it is to be merged. As illustrated, both the tracking maps received from portable devices and environment maps, serving as a set of canonical maps for localization and merging, are initially aligned with respect to gravity. Such an alignment, for example, may be achieved by aligning with gravity one coordinate of the coordinate frames used within those maps.

[0551] The map merge process 6900 begins with a map merge portion (e.g., map merge portion 810 of FIG. 26) receiving a tracking map (e.g., a new tracking map) from a portable device and one or more environment maps stored in a database in the cloud. Receiving the tracking map may include receiving data associated with the tracking map (e.g., sets of feature points and/or other information described herein), where the data is organized with respect to gravity. The tracking map may be represented by a coordinate system, where one of the coordinates of the coordinate system is the direction of gravity.

[0552] As described herein, the one or more environment maps may be selected from the database as candidates for comparison to the tracking map. This selection may be based on area identifiers, area attributes, or other metadata associated with the new tracking map. In some embodiments, the environment maps may be derived by filtering a larger

universe of canonical maps based on one or more criteria including, for example, a geographic location of the new tracking map, similarity of one or more identifiers of network access points associated with the new tracking map and the environment maps, similarity of metrics representing contents of the new tracking map and the environment maps, and degree of match between a portion of the new tracking map and portions of the environment maps. The one or more environment maps received by the map merge portion may include a filtered set of one or more environment maps.

[0553] In some embodiments, the one or more candidate environment maps may be ranked, for example by a map rank portion (e.g., map rank portion **806** of FIG. 26), for use in selecting the environment maps to merge with the new tracking map. The environment maps may be ranked by relevance. The process **6900** may be performed on the candidate maps one at a time, starting with the highest ranking map. If a suitable transformation between that map and the tracking map is found, the new tracking map may be merged with the selected candidate environment map to create one or more updated/merged environment maps as described herein. If no suitable transformation is found, the process **6900** may be repeated with the next highest ranked candidate map. In embodiments in which maps are segmented into areas, such as tiles, the process **6900** map operate on tiles or other portions of maps.

[0554] As shown in FIG. 69, at act **6910** the tracking map and the selected environment map are aligned. As described above, alignment may entail searching for a transformation between features in the new tracking map and the selected environment map that have been determined to correspond, such as a result of having matching identifiers. Following act **6910**, the determined transformation may be applied to the tracking map at act **6912** such that it is aligned with the selected environment map. At act **6920**, a determination is made as to whether the aligned tracking map transformed in act **6912** aligns with a gravity direction of the environment map. The process **6900** may branch at decision block **6922**, depending on whether the transformed tracking map has a valid alignment with respect to gravity. A valid alignment may be identified, for example, based on a rotation, less than a threshold amount, between the gravity direction of the tracking map and the gravity direction of the environment map.

[0555] In response to a determination that the gravity directions of the two maps align with one another, the maps may be merged in act **6930**. The resulting merged map retains its orientation with respect to gravity.

[0556] In response to a determination, in act **6920**, that the gravity directions of the two maps do not align with one another, or other indication that the gravity direction of the transformed tracking map is invalid, the map merge process refrains from merging the new tracking map with the environment map, in act **6930**. Rather, in response to a determination, in act **6920**, that the gravity direction of the transformed tracking map is invalid, process **6900** may branch from decision block **6922** back to act **6910** where further transformations are attempted to determine whether the gravity directions of the two maps can be aligned. The process **6900** may continue looping in this way until a valid alignment is found or some stop condition on the process is reached. That stop condition may be, for example, based on

the passage of time, attempting a predetermined number of transformations, or other conditions.

[0557] Upon reaching such a stop condition, one or more actions may be performed. In some embodiments, the map merge process **6900** may be performed for each environment map in the set of candidate environment maps until an environment map is located that can be aligned with the transformed tracking map while preserving the orientation of the tracking map with respect to gravity. The new tracking map may then be merged with the located environment map. Alternatively, if no such candidate map is identified, other actions may be taken, such as saving the tracking map as a new canonical map, or ending the merge process without updating the stored set of canonical maps based on the new tracking map.

[0558] FIG. 69 illustrates a merge process in which a tracking map is merged with an environment map selected from a database of canonical maps. Similar processing may be applied to other map merge operations. For example, as canonical maps in the database expand, such as through merging with new tracking maps, the canonical maps may grow such that two or more maps in the database represent overlapping regions of the physical world. In some embodiments, a process as illustrated in FIG. 69 may be applied with selected maps from the database as inputs in place of a tracking map so that overlapping canonical maps may be merged. Regardless of the specific maps that are input to the merge process, the process may be implemented so as to preserve the orientation of the maps with respect to gravity, such as by refraining from merging maps based on an alignment that changes the orientation with respect to gravity of the map to be merged.

Further Considerations

[0559] FIG. 60 shows a diagrammatic representation of a machine in the exemplary form of a computer system **1900** within which a set of instructions, for causing the machine to perform any one or more of the methodologies discussed herein, may be executed, according to some embodiments. In alternative embodiments, the machine operates as a standalone device or may be connected (e.g., networked) to other machines. Further, while only a single machine is illustrated, the term "machine" shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

[0560] The exemplary computer system **1900** includes a processor **1902** (e.g., a central processing unit (CPU), a graphics processing unit (GPU) or both), a main memory **1904** (e.g., read only memory (ROM), flash memory, dynamic random access memory (DRAM) such as synchronous DRAM (SDRAM) or Rambus DRAM (RDRAM), etc.), and a static memory **1906** (e.g., flash memory, static random access memory (SRAM), etc.), which communicate with each other via a bus **1908**.

[0561] The computer system **1900** may further include a disk drive unit **1916**, and a network interface device **1920**.

[0562] The disk drive unit **1916** includes a machine-readable medium **1922** on which is stored one or more sets of instructions **1924** (e.g., software) embodying any one or more of the methodologies or functions described herein. The software may also reside, completely or at least partially, within the main memory **1904** and/or within the processor **1902** during execution thereof by the computer

system **1900**, the main memory **1904** and the processor **1902** also constituting machine-readable media.

[0563] The software may further be transmitted or received over a network **18** via the network interface device **1920**.

[0564] The computer system **1900** includes a driver chip **1950** that is used to drive projectors to generate light. The driver chip **1950** includes its own data store **1960** and its own processor **1962**.

[0565] While the machine-readable medium **1922** is shown in an exemplary embodiment to be a single medium, the term "machine-readable medium" should be taken to include a single medium or multiple media (e.g., a centralized or distributed database, and/or associated caches and servers) that store the one or more sets of instructions. The term "machine-readable medium" shall also be taken to include any medium that is capable of storing, encoding, or carrying a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present invention. The term "machine-readable medium" shall accordingly be taken to include, but not be limited to, solid-state memories, optical and magnetic media, and carrier wave signals.

[0566] Having thus described several aspects of some embodiments, it is to be appreciated that various alterations, modifications, and improvements will readily occur to those skilled in the art.

[0567] As one example, embodiments are described in connection with an augmented (AR) environment. It should be appreciated that some or all of the techniques described herein may be applied in an MR environment or more generally in other XR environments, and in VR environments.

[0568] As another example, embodiments are described in connection with devices, such as wearable devices. It should be appreciated that some or all of the techniques described herein may be implemented via networks (such as cloud), discrete applications, and/or any suitable combinations of devices, networks, and discrete applications.

[0569] Further, FIG. 29 provides examples of criteria that may be used to filter candidate maps to yield a set of high ranking maps. Other criteria may be used instead of or in addition to the described criteria. For example, if multiple candidate maps have similar values of a metric used for filtering out less desirable maps, characteristics of the candidate maps may be used to determine which maps are retained as candidate maps or filtered out. For example, larger or more dense candidate maps may be prioritized over smaller candidate maps.

[0570] Such alterations, modifications, and improvements are intended to be part of this disclosure, and are intended to be within the spirit and scope of the disclosure. Further, though advantages of the present disclosure are indicated, it should be appreciated that not every embodiment of the disclosure will include every described advantage. Some embodiments may not implement any features described as advantageous herein and in some instances. Accordingly, the foregoing description and drawings are by way of example only.

[0571] The above-described embodiments of the present disclosure can be implemented in any of numerous ways. For example, the embodiments may be implemented using hardware, software or a combination thereof. When implemented in software, the software code can be executed on

any suitable processor or collection of processors, whether provided in a single computer or distributed among multiple computers. Such processors may be implemented as integrated circuits, with one or more processors in an integrated circuit component, including commercially available integrated circuit components known in the art by names such as CPU chips, GPU chips, microprocessor, microcontroller, or co-processor. In some embodiments, a processor may be implemented in custom circuitry, such as an ASIC, or semicustom circuitry resulting from configuring a programmable logic device. As yet a further alternative, a processor may be a portion of a larger circuit or semiconductor device, whether commercially available, semi-custom or custom. As a specific example, some commercially available microprocessors have multiple cores such that one or a subset of those cores may constitute a processor. Though, a processor may be implemented using circuitry in any suitable format.

[0572] Further, it should be appreciated that a computer may be embodied in any of a number of forms, such as a rack-mounted computer, a desktop computer, a laptop computer, or a tablet computer. Additionally, a computer may be embedded in a device not generally regarded as a computer but with suitable processing capabilities, including a Personal Digital Assistant (PDA), a smart phone or any other suitable portable or fixed electronic device.

[0573] Also, a computer may have one or more input and output devices. These devices can be used, among other things, to present a user interface. Examples of output devices that can be used to provide a user interface include printers or display screens for visual presentation of output and speakers or other sound generating devices for audible presentation of output. Examples of input devices that can be used for a user interface include keyboards, and pointing devices, such as mice, touch pads, and digitizing tablets. As another example, a computer may receive input information through speech recognition or in other audible format. In the embodiment illustrated, the input/output devices are illustrated as physically separate from the computing device. In some embodiments, however, the input and/or output devices may be physically integrated into the same unit as the processor or other elements of the computing device. For example, a keyboard might be implemented as a soft keyboard on a touch screen. In some embodiments, the input/output devices may be entirely disconnected from the computing device, and functionally integrated through a wireless connection.

[0574] Such computers may be interconnected by one or more networks in any suitable form, including as a local area network or a wide area network, such as an enterprise network or the Internet. Such networks may be based on any suitable technology and may operate according to any suitable protocol and may include wireless networks, wired networks or fiber optic networks.

[0575] Also, the various methods or processes outlined herein may be coded as software that is executable on one or more processors that employ any one of a variety of operating systems or platforms. Additionally, such software may be written using any of a number of suitable programming languages and/or programming or scripting tools, and also may be compiled as executable machine language code or intermediate code that is executed on a framework or virtual machine.

[0576] In this respect, the disclosure may be embodied as a computer readable storage medium (or multiple computer

readable media) (e.g., a computer memory, one or more floppy discs, compact discs (CD), optical discs, digital video disks (DVD), magnetic tapes, flash memories, circuit configurations in Field Programmable Gate Arrays or other semiconductor devices, or other tangible computer storage medium) encoded with one or more programs that, when executed on one or more computers or other processors, perform methods that implement the various embodiments of the disclosure discussed above. As is apparent from the foregoing examples, a computer readable storage medium may retain information for a sufficient time to provide computer-executable instructions in a non-transitory form. Such a computer readable storage medium or media can be transportable, such that the program or programs stored thereon can be loaded onto one or more different computers or other processors to implement various aspects of the present disclosure as discussed above. As used herein, the term "computer-readable storage medium" encompasses only a computer-readable medium that can be considered to be a manufacture (i.e., article of manufacture) or a machine. In some embodiments, the disclosure may be embodied as a computer readable medium other than a computer-readable storage medium, such as a propagating signal.

[0577] The terms "program" or "software" are used herein in a generic sense to refer to any type of computer code or set of computer-executable instructions that can be employed to program a computer or other processor to implement various aspects of the present disclosure as discussed above. Additionally, it should be appreciated that according to one aspect of this embodiment, one or more computer programs that when executed perform methods of the present disclosure need not reside on a single computer or processor, but may be distributed in a modular fashion amongst a number of different computers or processors to implement various aspects of the present disclosure.

[0578] Computer-executable instructions may be in many forms, such as program modules, executed by one or more computers or other devices. Generally, program modules include routines, programs, objects, components, data structures, etc. that perform particular tasks or implement particular abstract data types. Typically the functionality of the program modules may be combined or distributed as desired in various embodiments.

[0579] Also, data structures may be stored in computer-readable media in any suitable form. For simplicity of illustration, data structures may be shown to have fields that are related through location in the data structure. Such relationships may likewise be achieved by assigning storage for the fields with locations in a computer-readable medium that conveys relationship between the fields. However, any suitable mechanism may be used to establish a relationship between information in fields of a data structure, including through the use of pointers, tags or other mechanisms that establish relationship between data elements.

[0580] Various aspects of the present disclosure may be used alone, in combination, or in a variety of arrangements not specifically discussed in the embodiments described in the foregoing and is therefore not limited in its application to the details and arrangement of components set forth in the foregoing description or illustrated in the drawings. For example, aspects described in one embodiment may be combined in any manner with aspects described in other embodiments.

[0581] Also, the disclosure may be embodied as a method, of which an example has been provided. The acts performed as part of the method may be ordered in any suitable way. Accordingly, embodiments may be constructed in which acts are performed in an order different than illustrated, which may include performing some acts simultaneously, even though shown as sequential acts in illustrative embodiments.

[0582] Use of ordinal terms such as "first," "second," "third," etc., in the claims to modify a claim element does not by itself connote any priority, precedence, or order of one claim element over another or the temporal order in which acts of a method are performed, but are used merely as labels to distinguish one claim element having a certain name from another element having a same name (but for use of the ordinal term) to distinguish the claim elements.

[0583] Also, the phraseology and terminology used herein is for the purpose of description and should not be regarded as limiting. The use of "including," "comprising," or "having," "containing," "involving," and variations thereof herein, is meant to encompass the items listed thereafter and equivalents thereof as well as additional items.

1. A method of operating a cross reality system in which environment maps are stored in a database, the method comprising:

determining whether to merge a first map and a second map, wherein the first map and the second map are aligned with respect to a gravity direction, and the determining comprises:

applying a plurality of transformations to at least a portion of the first map, wherein each of the plurality of applied transformations preserves an orientation of the first map with respect to gravity;

for each applied transformation of the plurality of transformations, computing an error in alignment of the first map and the second map; and

selecting an applied transformation of the plurality of transformations based on the applied transformation having a low error relative to other applied transformation of the plurality of transformations such that the determining is based on the error of the selected applied transformation.

2. The method of claim 1, wherein the transformation applied to at least the portion of the first map to align it with the second map is based at least in part on matching sets of features.

3. The method of claim 1, wherein at least one of the first map and the second map is accessed from the database.

4. The method of claim 1, wherein at least one of the first map and the second map is constructed from information collected by at least one user device.

5. The method of claim 1, further comprising storing a canonical map resulting from merging the first map and the second map in the database.

6. The method of claim 1, wherein at least one of the first map and the second map is a canonical map resulting from a previous merge process.

7. The method of claim 1, wherein the first map and the second map represent overlapping portions of a physical world.

8. The method of claim 1, wherein the first map and the second map represent portions of a physical world that do not overlap.

9. The method of claim **1**, further comprising, based on a determination that the applied transformation does not have the low error relative to the other applied transformation of the plurality of transformation, identifying a further transformation and repeating the determining whether a transformation results in alignment with respect to the gravity direction.

10. The method of claim **1**, wherein determining whether a result in alignment with respect to the gravity direction comprises comparing the error to a threshold.

11. A computing device configured for use in a cross reality system, the computing device comprising:
at least one processor;
a computer-readable medium connected to the processor;
a plurality of environment maps stored in the computer-readable medium; and
computer-executable instructions configured to, when executed by the at least one processor, perform a method comprising:
determining whether to merge a first map and a second map, wherein the first map and the second map are aligned with respect to a gravity direction, and the determining comprises:
applying a plurality of transformations to at least a portion of the first map, wherein each of the plurality of applied transformations preserves an orientation of the first map with respect to gravity;
for each applied transformation of the plurality of transformations, computing an error in alignment of the first map and the second map; and
selecting an applied transformation of the plurality of transformations based on the applied transformation having a low error relative to other applied transformation of the plurality of transformations such that the determining is based on the error of the selected applied transformation.

12. The computing device of claim **11**, wherein the method further comprises storing a canonical map resulting from merging the first map and the second map in the computer-readable medium.

13. The computing device of claim **11**, wherein at least one of the first map and the second map is constructed from information collected by at least one user device.

14. The computing device of claim **11**, further comprising, based on a determination that the applied transformation does not have the low error relative to the other applied transformation of the plurality of transformation, identifying a further transformation and repeating the determining whether a transformation results in alignment with respect to the gravity direction.

15. The computing device of claim **11**, wherein determining whether a result in alignment with respect to the gravity direction comprises comparing the error to a threshold.

16. A cloud computing environment for an augmented reality system configured for communication with a plurality of user devices comprising sensors, comprising:

a database storing a plurality of environment maps constructed from data supplied by the plurality of user devices; and
non-transitory computer storage media storing computer-executable instructions that, when executed by at least one processor in the cloud computing environment, perform a method comprising:
determining whether to merge a first map and a second map, wherein the first map and the second map are aligned with respect to a gravity direction, and the determining comprises:
applying a plurality of transformations to at least a portion of the first map, wherein each of the plurality of applied transformations preserves an orientation of the first map with respect to gravity;
for each applied transformation of the plurality of transformations, computing an error in alignment of the first map and the second map; and
selecting an applied transformation of the plurality of transformations based on the applied transformation having a low error relative to other applied transformation of the plurality of transformations such that the determining is based on the error of the selected applied transformation.

17. The cloud computing environment of claim **16**, the method further comprising storing a canonical map resulting from merging the first environment map and the second environment map in the database.

18. The cloud computing environment of claim **16**, wherein at least one of the first map and the second map is constructed from information collected by at least one user device.

19. The cloud computing environment of claim **16**, further comprising, based on a determination that the applied transformation does not have the low error relative to the other applied transformation of the plurality of transformation, identifying a further transformation and repeating the determining whether a transformation results in alignment with respect to the gravity direction.

20. The cloud computing environment of claim **16**, wherein determining whether a result in alignment with respect to the gravity direction comprises comparing the error to a threshold.

* * * * *