



US 20250259629A1

(19) United States

(12) Patent Application Publication

WALKER, II et al.

(10) Pub. No.: US 2025/0259629 A1

(43) Pub. Date: Aug. 14, 2025

(54) APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT

(71) Applicant: Apple Inc., Cupertino, CA (US)

(72) Inventors: Robert A. WALKER, II, San Francisco, CA (US); Brandon J. NEWENDORP, San Francisco, CA (US); Rohit DASARI, San Francisco, CA (US); Richard D. GIULI, Arroyo Grande, CA (US); Thomas R. GRUBER, Santa Cruz, CA (US); Carey E. RADEBAUGH, Los Gatos, CA (US); Ashish GARG, Bellevue, CA (US); Vineet KHOSLA, Pacifica, CA (US); Jonathan H. RUSSELL, Incline Village, NV (US); Corey PETERSON, Saratoga, CA (US)

(21) Appl. No.: 19/193,579

(22) Filed: Apr. 29, 2025

Related U.S. Application Data

(63) Continuation of application No. 18/226,765, filed on Jul. 26, 2023, now Pat. No. 12,293,763, which is a continuation of application No. 17/497,734, filed on Oct. 8, 2021, now Pat. No. 11,749,275, which is a continuation of application No. 16/737,100, filed on Jan. 8, 2020, now Pat. No. 11,152,002, which is a continuation of application No. 16/207,901, filed on Dec. 3, 2018, now Pat. No. 10,580,409, which is a continuation of application No. 15/269,728, filed on Sep. 19, 2016, now Pat. No. 10,297,253.

(60) Provisional application No. 62/348,929, filed on Jun. 11, 2016.

Publication Classification

(51) Int. Cl.


G10L 15/22 (2006.01)
G06F 40/35 (2020.01)
G10L 13/02 (2013.01)
G10L 15/06 (2013.01)
G10L 15/18 (2013.01)
G10L 15/30 (2013.01)

(52) U.S. Cl.

CPC *G10L 15/22* (2013.01); *G10L 13/02* (2013.01); *G10L 15/1815* (2013.01); *G06F 40/35* (2020.01); *G10L 15/063* (2013.01); *G10L 2015/223* (2013.01); *G10L 2015/228* (2013.01); *G10L 15/30* (2013.01)

(57) ABSTRACT

Systems and processes for application integration with a digital assistant are provided. In accordance with one example, a method includes, at an electronic device having one or more processors and memory, receiving a natural-language user input; identifying, with the one or more processors, an intent object of a set of intent objects and a parameter associated with the intent, where the intent object and the parameter are derived from the natural-language user input. The method further includes identifying a software application associated with the intent object of the set of intent objects; and providing the intent object and the parameter to the software application.







FIG. 1

**FIG. 2A**




FIG. 2B





FIG. 3

**FIG. 4**




FIG. 5A





FIG. 5B




FIG. 6A




FIG. 6B




FIG. 7A




FIG. 7B




FIG. 7C

800




FIG. 8

Attorney Docket No.: P30368USC5/77670000248205
14/21

900




FIG. 9




FIG. 10C




FIG. 10B




FIG. 10A




FIG. 10D




FIG. 10E




FIG. 11




FIG. 12

1300 ↗




FIG. 13




FIG. 14

APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT

CROSS-REFERENCE TO RELATED APPLICATIONS

[0001] This application is a continuation of U.S. patent application Ser. No. 18/226,765, entitled “APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT”, filed on Jul. 26, 2023, which is a continuation of U.S. patent application Ser. No. 17/497,734, now U.S. Pat. No. 11,749,275, entitled “APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT”, filed on Oct. 8, 2021, which is a continuation of U.S. patent application Ser. No. 16/737,100, now U.S. Pat. No. 11,152,002, entitled “APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT”, filed on Jan. 8, 2020, which is a continuation of U.S. patent application Ser. No. 16/207,901, now U.S. Pat. No. 10,580,409, entitled “APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT”, filed on Dec. 3, 2018, which is a continuation of U.S. patent application Ser. No. 15/269,728, now U.S. Pat. No. 10,297,253, entitled “APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT,” filed on Sep. 19, 2016, which claims priority to U.S. Provisional Patent Application Ser. No. 62/348,929, entitled “APPLICATION INTEGRATION WITH A DIGITAL ASSISTANT,” filed on Jun. 11, 2016, all of which are hereby incorporated by reference in their entirety for all purposes.

FIELD

[0002] The present disclosure relates generally to interfacing with applications, and more specifically to techniques for application integration with a digital assistant.

BACKGROUND

[0003] Digital assistants can help a user perform various functions on a user device. For example, digital assistants can set alarms, provide weather updates, and perform searches both locally and on the Internet, all while providing a natural-language interface to a user. Existing digital assistants, however, cannot effectively integrate with applications, such as those locally stored on the user device, and in particular, third-party applications. Accordingly, existing digital assistants cannot provide a natural-language interface with such applications.

BRIEF SUMMARY

[0004] Example methods are disclosed herein. An example method includes, at an electronic device having one or more processors, receiving a natural-language user input, and identifying, with the one or more processors, an intent of a set of intents and a parameter associated with the intent, where the intent and the parameter are derived from the natural-language user input. The method further includes identifying a software application associated with the intent of the set of intents and providing the intent and the parameter to the software application.

[0005] An example method includes, at one or more electronic devices, each having one or more processors, receiving a natural-language user input; determining, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; identify-

ing a software application based on at least one of the intent or the parameter; and providing the intent and the parameter to the software application.

[0006] An example method includes, at one or more electronic devices, each having one or more processors, receiving a natural-language user input; identifying, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; and determining, based on at least one of the intent or the parameter, whether a task corresponding to the intent can be fulfilled. The method further includes, in accordance with a determination that the task corresponding to the intent can be fulfilled, providing the intent and the parameter to a software application associated with the intent, and, in accordance with a determination that the task corresponding to the intent cannot be fulfilled, providing a list of one or more software applications associated with the intent.

[0007] An example method includes, at a first electronic device having one or more processors, receiving a natural-language user input, where the natural-language user input is indicative of an intent of a set of intents; providing the natural-language user input to a second electronic device; and receiving, from the second electronic device, an indication that a software application associated with the intent is not located on the first electronic device. The method further includes, responsive to the notification, obtaining the list of applications associated with the intent; displaying, with a touch-sensitive display of the first electronic device, the list of applications associated with the intent in a user interface; receiving a user input indicative of a selection of an application of the list of applications; and providing the intent of the set of intents to the application.

[0008] Example non-transitory computer-readable media are disclosed herein. An example non-transitory computer-readable storage medium stores one or more programs. The one or more programs comprise instructions, which when executed by one or more processors of an electronic device, cause the electronic device to receive a natural-language user input; identify, with the one or more processors, an intent of a set of intents and a parameter associated with the intent, wherein the intent and the parameter are derived from the natural-language user input; identify a software application associated with the intent of the set of intents; and provide the intent and the parameter to the software application.

[0009] Example non-transitory computer-readable storage media store one or more programs. The one or more programs comprise instructions, which when executed by one or more processors of one or more electronic devices, cause the one or more electronic devices to receive a natural-language user input; determine, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; identify a software application based on at least one of the intent or the parameter; and provide the intent and the parameter to the software application.

[0010] Example non-transitory computer-readable storage media store one or more programs. The one or more programs comprise instructions, which when executed by one or more processors of one or more electronic devices, cause the one or more electronic devices to receive a natural-language user input; identify, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; determine, based on at least one

of the intent or the parameter, whether a task corresponding to the intent can be fulfilled; in accordance with a determination that the task corresponding to the intent can be fulfilled, provide the intent and the parameter to a software application associated with the intent; and in accordance with a determination that the task corresponding to the intent cannot be fulfilled, provide a list of one or more software applications associated with the intent.

[0011] An example non-transitory computer-readable storage medium stores one or more programs. The one or more programs comprise instructions, which when executed by one or more processors of an electronic device, cause the electronic device to receive a natural-language user input, wherein the natural-language user input is indicative of an intent of a set of intents; provide the natural-language user input to a second electronic device; receive, from the second electronic device, an indication that a software application associated with the intent is not located on the first electronic device; responsive to the notification, obtain the list of applications associated with the intent; display, with a touch-sensitive display of the first electronic device, the list of applications associated with the intent in a user interface; receive a user input indicative of a selection of an application of the list of applications; and provide the intent of the set of intents to the application.

[0012] Example electronic devices and systems are disclosed herein. An example electronic device comprises one or more processors; a memory; and one or more programs, where the one or more programs are stored in the memory and configured to be executed by the one or more processors, the one or more programs including instructions for receiving a natural-language user input; identifying, with the one or more processors, an intent of a set of intents and a parameter associated with the intent, wherein the intent and the parameter are derived from the natural-language user input; identifying a software application associated with the intent of the set of intents; and providing the intent and the parameter to the software application.

[0013] An example system includes one or more processors of one or more electronic devices; one or more memories of the one or more electronic devices; and one or more programs, where the one or more programs are stored in the one or more memories and configured to be executed by the one or more processors, the one or more programs including instructions for receiving a natural-language user input; determining, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; identifying a software application based on at least one of the intent or the parameter; and providing the intent and the parameter to the software application.

[0014] An example system includes one or more processors of one or more electronic devices; one or more memories of the one or more electronic devices; and one or more programs, where the one or more programs are stored in the one or more memories and configured to be executed by the one or more processors, the one or more programs including instructions for receiving a natural-language user input; identifying, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; determining, based on at least one of the intent or the parameter, whether a task corresponding to the intent can be fulfilled; in accordance with a determination that the task corresponding to the intent can be fulfilled, providing the intent and the parameter to a software application associated

with the intent; and in accordance with a determination that the task corresponding to the intent cannot be fulfilled, providing a list of one or more software applications associated with the intent.

[0015] An example electronic device comprises one or more processors; a memory; and one or more programs, where the one or more programs are stored in the memory and configured to be executed by the one or more processors, the one or more programs including instructions for receiving a natural-language user input, wherein the natural-language user input is indicative of an intent of a set of intents; providing the natural-language user input to a second electronic device; receiving, from the second electronic device, an indication that a software application associated with the intent is not located on the first electronic device; responsive to the notification, obtaining the list of applications associated with the intent; displaying, with a touch-sensitive display of the first electronic device, the list of applications associated with the intent in a user interface; receiving a user input indicative of a selection of an application of the list of applications; and providing the intent of the set of intents to the application.

[0016] An example electronic device comprises means for receiving a natural-language user input; means for identifying an intent of a set of intents and a parameter associated with the intent, where the intent and the parameter are derived from the natural-language user input; means for identifying a software application associated with the intent of the set of intents; and means for providing the intent and the parameter to the software application.

[0017] An example system comprises means for receiving a natural-language user input; means for determining, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; means for identifying a software application based on at least one of the intent or the parameter; and means for providing the intent and the parameter to the software application.

[0018] An example system comprises means for receiving a natural-language user input; means for identifying, based on the natural-language user input, an intent of a set of intents and a parameter associated with the intent; means for determining, based on at least one of the intent or the parameter, whether a task corresponding to the intent can be fulfilled; means for, in accordance with a determination that the task corresponding to the intent can be fulfilled, providing the intent and the parameter to a software application associated with the intent; and means for, in accordance with a determination that the task corresponding to the intent cannot be fulfilled, providing a list of one or more software applications associated with the intent.

[0019] An example electronic device includes means for receiving a natural-language user input, wherein the natural-language user input is indicative of an intent of a set of intents; means for providing the natural-language user input to a second electronic device; means for receiving, from the second electronic device, an indication that a software application associated with the intent is not located on the first electronic device; means for, responsive to the notification, obtaining the list of applications associated with the intent; means for displaying, with a touch-sensitive display of the first electronic device, the list of applications associated with the intent in a user interface; means for receiving a user input indicative of a selection of an application of the

list of applications; and means for providing the intent of the set of intents to the application.

BRIEF DESCRIPTION OF THE FIGURES

[0020] For a better understanding of the various described embodiments, reference should be made to the Detailed Description below, in conjunction with the following drawings in which like reference numerals refer to corresponding parts throughout the figures.

[0021] FIG. 1 is a block diagram illustrating a system and environment for implementing a digital assistant according to various examples.

[0022] FIG. 2A is a block diagram illustrating a portable multifunction device implementing the client-side portion of a digital assistant in accordance with some embodiments.

[0023] FIG. 2B is a block diagram illustrating exemplary components for event handling according to various examples.

[0024] FIG. 3 illustrates a portable multifunction device implementing the client-side portion of a digital assistant according to various examples.

[0025] FIG. 4 is a block diagram of an exemplary multifunction device with a display and a touch-sensitive surface according to various examples.

[0026] FIG. 5A illustrates an exemplary user interface for a menu of applications on a portable multifunction device according to various examples.

[0027] FIG. 5B illustrates an exemplary user interface for a multifunction device with a touch-sensitive surface that is separate from the display according to various examples.

[0028] FIG. 6A illustrates a personal electronic device according to various examples.

[0029] FIG. 6B is a block diagram illustrating a personal electronic device according to various examples.

[0030] FIG. 7A is a block diagram illustrating a digital assistant system or a server portion thereof according to various examples.

[0031] FIG. 7B illustrates the functions of the digital assistant shown in FIG. 7A according to various examples.

[0032] FIG. 7C illustrates a portion of an ontology according to various examples.

[0033] FIG. 8 illustrates a flow diagram of a process for operating a digital assistant in accordance with some embodiments.

[0034] FIG. 9 illustrates a flow diagram of a process for operating a digital assistant in accordance with some embodiments.

[0035] FIGS. 10A-10C illustrate exemplary user interfaces of an electronic device in accordance with some embodiments.

[0036] FIGS. 10D-10E illustrate exemplary data flows of a digital assistant system in accordance with some embodiments.

[0037] FIGS. 11-14 illustrate functional block diagram of electronic devices in accordance with some embodiments.

DETAILED DESCRIPTION

[0038] In the following description of the disclosure and embodiments, reference is made to the accompanying drawings in which it is shown by way of illustration of specific embodiments that can be practiced. It is to be understood

that other embodiments and examples can be practiced and changes can be made without departing from the scope of the disclosure.

[0039] Although the following description uses terms “first,” “second,” etc. to describe various elements, these elements should not be limited by the terms. These terms are only used to distinguish one element from another. For example, a first input could be termed a second input, and, similarly, a second input could be termed a first input, without departing from the scope of the various described examples. The first input and the second input can both be outputs and, in some cases, can be separate and different inputs.

[0040] The terminology used in the description of the various described examples herein is for the purpose of describing particular examples only and is not intended to be limiting. As used in the description of the various described examples and the appended claims, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will also be understood that the term “and/or” as used herein refers to and encompasses any and all possible combinations of one or more of the associated listed items. It will be further understood that the terms “includes,” “including,” “comprises,” and/or “comprising,” when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof.

[0041] The term “if” may be construed to mean “when” or “upon” or “in response to determining” or “in response to detecting,” depending on the context. Similarly, the phrase “if it is determined” or “if [a stated condition or event] is detected” may be construed to mean “upon determining” or “in response to determining” or “upon detecting [the stated condition or event]” or “in response to detecting [the stated condition or event],” depending on the context.

1. System and Environment

[0042] FIG. 1 illustrates a block diagram of system 100 according to various examples. In some examples, system 100 can implement a digital assistant. The terms “digital assistant,” “virtual assistant,” “intelligent automated assistant,” or “automatic digital assistant” can refer to any information processing system that interprets natural language input in spoken and/or textual form to infer user intent, and performs actions (e.g., tasks) based on the inferred user intent. For example, to act on an inferred user intent, the system can perform one or more of the following: identifying a task flow with steps and parameters designed to accomplish the inferred user intent, inputting specific requirements from the inferred user intent into the task flow; executing the task flow by invoking programs, methods, services, APIs, or the like; and generating output responses to the user in an audible (e.g., speech) and/or visual form.

[0043] Specifically, a digital assistant can be capable of accepting a user request at least partially in the form of a natural language command, request, statement, narrative, and/or inquiry. Typically, the user request can seek either an informational answer or performance of a task by the digital assistant. A satisfactory response to the user request can be a provision of the requested informational answer, a performance of the requested task, or a combination of the two.

For example, a user can ask the digital assistant a question, such as "Where am I right now?" Based on the user's current location, the digital assistant can answer, "You are in Central Park near the west gate." The user can also request the performance of a task, for example, "Please invite my friends to my girlfriend's birthday party next week." In response, the digital assistant can acknowledge the request by saying "Yes, right away," and then send a suitable calendar invite on behalf of the user to each of the user's friends listed in the user's electronic address book. During performance of a requested task, the digital assistant can sometimes interact with the user in a continuous dialogue involving multiple exchanges of information over an extended period of time. There are numerous other ways of interacting with a digital assistant to request information or performance of various tasks. In addition to providing verbal responses and taking programmed actions, the digital assistant can also provide responses in other visual or audio forms, e.g., as text, alerts, music, videos, animations, etc.

[0044] As shown in FIG. 1, in some examples, a digital assistant can be implemented according to a client-server model. The digital assistant can include client-side portion 102 (hereafter "DA client 102") executed on user device 104 and server-side portion 106 (hereafter "DA server 106") executed on server system 108. DA client 102 can communicate with DA server 106 through one or more networks 110. DA client 102 can provide client-side functionalities such as user-facing input and output processing and communication with DA server 106. DA server 106 can provide server-side functionalities for any number of DA clients 102 each residing on a respective user device 104.

[0045] In some examples, DA server 106 can include client-facing I/O interface 112, one or more processing modules 114, data and models 116, and I/O interface to external services 118. The client-facing I/O interface 112 can facilitate the client-facing input and output processing for DA server 106. One or more processing modules 114 can utilize data and models 116 to process speech input and determine the user's intent based on natural language input. Further, one or more processing modules 114 perform task execution based on inferred user intent. In some examples, DA server 106 can communicate with external services 120 through network(s) 110 for task completion or information acquisition. I/O interface to external services 118 can facilitate such communications.

[0046] User device 104 can be any suitable electronic device. For example, user devices can be a portable multifunctional device (e.g., device 200, described below with reference to FIG. 2A), a multifunctional device (e.g., device 400, described below with reference to FIG. 4), or a personal electronic device (e.g., device 600, described below with reference to FIG. 6A-6B.) A portable multifunctional device can be, for example, a mobile telephone that also contains other functions, such as PDA and/or music player functions. Specific examples of portable multifunction devices can include the iPhone®, iPod Touch®, and iPad® devices from Apple Inc. of Cupertino, California. Other examples of portable multifunction devices can include, without limitation, laptop or tablet computers. Further, in some examples, user device 104 can be a non-portable multifunctional device. In particular, user device 104 can be a desktop computer, a game console, a television, or a television set-top box. In some examples, user device 104 can include a touch-sensitive surface (e.g., touch screen displays and/or

touchpads). Further, user device 104 can optionally include one or more other physical user-interface devices, such as a physical keyboard, a mouse, and/or a joystick. Various examples of electronic devices, such as multifunctional devices, are described below in greater detail.

[0047] Examples of communication network(s) 110 can include local area networks (LAN) and wide area networks (WAN), e.g., the Internet. Communication network(s) 110 can be implemented using any known network protocol, including various wired or wireless protocols, such as, for example, Ethernet, Universal Serial Bus (USB), FIREWIRE, Global System for Mobile Communications (GSM), Enhanced Data GSM Environment (EDGE), code division multiple access (CDMA), time division multiple access (TDMA), Bluetooth, Wi-Fi, voice over Internet Protocol (VoIP), Wi-MAX, or any other suitable communication protocol.

[0048] Server system 108 can be implemented on one or more standalone data processing apparatus or a distributed network of computers. In some examples, server system 108 can also employ various virtual devices and/or services of third-party service providers (e.g., third-party cloud service providers) to provide the underlying computing resources and/or infrastructure resources of server system 108.

[0049] In some examples, user device 104 can communicate with DA server 106 via second user device 122. Second user device 122 can be similar or identical to user device 104. For example, second user device 122 can be similar to devices 200, 400, or 600 described below with reference to FIGS. 2A, 4, and 6A-6B. User device 104 can be configured to communicatively couple to second user device 122 via a direct communication connection, such as Bluetooth, NFC, BTLE, or the like, or via a wired or wireless network, such as a local Wi-Fi network. In some examples, second user device 122 can be configured to act as a proxy between user device 104 and DA server 106. For example, DA client 102 of user device 104 can be configured to transmit information (e.g., a user request received at user device 104) to DA server 106 via second user device 122. DA server 106 can process the information and return relevant data (e.g., data content responsive to the user request) to user device 104 via second user device 122.

[0050] In some examples, user device 104 can be configured to communicate abbreviated requests for data to second user device 122 to reduce the amount of information transmitted from user device 104. Second user device 122 can be configured to determine supplemental information to add to the abbreviated request to generate a complete request to transmit to DA server 106. This system architecture can advantageously allow user device 104 having limited communication capabilities and/or limited battery power (e.g., a watch or a similar compact electronic device) to access services provided by DA server 106 by using second user device 122, having greater communication capabilities and/or battery power (e.g., a mobile phone, laptop computer, tablet computer, or the like), as a proxy to DA server 106. While only two user devices 104 and 122 are shown in FIG. 1, it should be appreciated that system 100 can include any number and type of user devices configured in this proxy configuration to communicate with DA server system 106.

[0051] Although the digital assistant shown in FIG. 1 can include both a client-side portion (e.g., DA client 102) and a server-side portion (e.g., DA server 106), in some examples, the functions of a digital assistant can be imple-

mented as a standalone application installed on a user device. In addition, the divisions of functionalities between the client and server portions of the digital assistant can vary in different implementations. For instance, in some examples, the DA client can be a thin-client that provides only user-facing input and output processing functions, and delegates all other functionalities of the digital assistant to a backend server.

2. Electronic Devices

[0052] Attention is now directed toward embodiments of electronic devices for implementing the client-side portion of a digital assistant. FIG. 2A is a block diagram illustrating portable multifunction device 200 with touch-sensitive display system 212 in accordance with some embodiments. Touch-sensitive display 212 is sometimes called a “touch screen” for convenience and is sometimes known as or called a “touch-sensitive display system.” Device 200 includes memory 202 (which optionally includes one or more computer-readable storage mediums), memory controller 222, one or more processing units (CPUs) 220, peripherals interface 218, RF circuitry 208, audio circuitry 210, speaker 211, microphone 213, input/output (I/O) subsystem 206, other input control devices 216, and external port 224. Device 200 optionally includes one or more optical sensors 264. Device 200 optionally includes one or more contact intensity sensors 265 for detecting intensity of contacts on device 200 (e.g., a touch-sensitive surface such as touch-sensitive display system 212 of device 200). Device 200 optionally includes one or more tactile output generators 267 for generating tactile outputs on device 200 (e.g., generating tactile outputs on a touch-sensitive surface such as touch-sensitive display system 212 of device 200 or touchpad 455 of device 400). These components optionally communicate over one or more communication buses or signal lines 203.

[0053] As used in the specification and claims, the term “intensity” of a contact on a touch-sensitive surface refers to the force or pressure (force per unit area) of a contact (e.g., a finger contact) on the touch-sensitive surface, or to a substitute (proxy) for the force or pressure of a contact on the touch-sensitive surface. The intensity of a contact has a range of values that includes at least four distinct values and more typically includes hundreds of distinct values (e.g., at least 256). Intensity of a contact is, optionally, determined (or measured) using various approaches and various sensors or combinations of sensors. For example, one or more force sensors underneath or adjacent to the touch-sensitive surface are, optionally, used to measure force at various points on the touch-sensitive surface. In some implementations, force measurements from multiple force sensors are combined (e.g., a weighted average) to determine an estimated force of a contact. Similarly, a pressure-sensitive tip of a stylus is, optionally, used to determine a pressure of the stylus on the touch-sensitive surface. Alternatively, the size of the contact area detected on the touch-sensitive surface and/or changes thereto, the capacitance of the touch-sensitive surface proximate to the contact and/or changes thereto, and/or the resistance of the touch-sensitive surface proximate to the contact and/or changes thereto are, optionally, used as a substitute for the force or pressure of the contact on the touch-sensitive surface. In some implementations, the substitute measurements for contact force or pressure are used directly to determine whether an intensity threshold has been

exceeded (e.g., the intensity threshold is described in units corresponding to the substitute measurements). In some implementations, the substitute measurements for contact force or pressure are converted to an estimated force or pressure, and the estimated force or pressure is used to determine whether an intensity threshold has been exceeded (e.g., the intensity threshold is a pressure threshold measured in units of pressure). Using the intensity of a contact as an attribute of a user input allows for user access to additional device functionality that may otherwise not be accessible by the user on a reduced-size device with limited real estate for displaying affordances (e.g., on a touch-sensitive display) and/or receiving user input (e.g., via a touch-sensitive display, a touch-sensitive surface, or a physical/mechanical control such as a knob or a button).

[0054] As used in the specification and claims, the term “tactile output” refers to physical displacement of a device relative to a previous position of the device, physical displacement of a component (e.g., a touch-sensitive surface) of a device relative to another component (e.g., housing) of the device, or displacement of the component relative to a center of mass of the device that will be detected by a user with the user’s sense of touch. For example, in situations where the device or the component of the device is in contact with a surface of a user that is sensitive to touch (e.g., a finger, palm, or other part of a user’s hand), the tactile output generated by the physical displacement will be interpreted by the user as a tactile sensation corresponding to a perceived change in physical characteristics of the device or the component of the device. For example, movement of a touch-sensitive surface (e.g., a touch-sensitive display or trackpad) is, optionally, interpreted by the user as a “down click” or “up click” of a physical actuator button. In some cases, a user will feel a tactile sensation such as an “down click” or “up click” even when there is no movement of a physical actuator button associated with the touch-sensitive surface that is physically pressed (e.g., displaced) by the user’s movements. As another example, movement of the touch-sensitive surface is, optionally, interpreted or sensed by the user as “roughness” of the touch-sensitive surface, even when there is no change in smoothness of the touch-sensitive surface. While such interpretations of touch by a user will be subject to the individualized sensory perceptions of the user, there are many sensory perceptions of touch that are common to a large majority of users. Thus, when a tactile output is described as corresponding to a particular sensory perception of a user (e.g., an “up click,” a “down click,” “roughness”), unless otherwise stated, the generated tactile output corresponds to physical displacement of the device or a component thereof that will generate the described sensory perception for a typical (or average) user.

[0055] It should be appreciated that device 200 is only one example of a portable multifunction device, and that device 200 optionally has more or fewer components than shown, optionally combines two or more components, or optionally has a different configuration or arrangement of the components. The various components shown in FIG. 2A are implemented in hardware, software, or a combination of both hardware and software, including one or more signal processing and/or application-specific integrated circuits.

[0056] Memory 202 may include one or more computer-readable storage mediums. The computer-readable storage mediums may be tangible and non-transitory. Memory 202 may include high-speed random access memory and may

also include non-volatile memory, such as one or more magnetic disk storage devices, flash memory devices, or other non-volatile solid-state memory devices. Memory controller 222 may control access to memory 202 by other components of device 200.

[0057] In some examples, a non-transitory computer-readable storage medium of memory 202 can be used to store instructions (e.g., for performing aspects of process 1100, described below) for use by or in connection with an instruction execution system, apparatus, or device, such as a computer-based system, processor-containing system, or other system that can fetch the instructions from the instruction execution system, apparatus, or device and execute the instructions. In other examples, the instructions (e.g., for performing aspects of process 1100, described below) can be stored on a non-transitory computer-readable storage medium (not shown) of the server system 108 or can be divided between the non-transitory computer-readable storage medium of memory 202 and the non-transitory computer-readable storage medium of server system 108. In the context of this document, a “non-transitory computer-readable storage medium” can be any medium that can contain or store the program for use by or in connection with the instruction execution system, apparatus, or device.

[0058] Peripherals interface 218 can be used to couple input and output peripherals of the device to CPU 220 and memory 202. The one or more processors 220 run or execute various software programs and/or sets of instructions stored in memory 202 to perform various functions for device 200 and to process data. In some embodiments, peripherals interface 218, CPU 220, and memory controller 222 may be implemented on a single chip, such as chip 204. In some other embodiments, they may be implemented on separate chips.

[0059] RF (radio frequency) circuitry 208 receives and sends RF signals, also called electromagnetic signals. RF circuitry 208 converts electrical signals to/from electromagnetic signals and communicates with communications networks and other communications devices via the electromagnetic signals. RF circuitry 208 optionally includes well-known circuitry for performing these functions, including but not limited to an antenna system, an RF transceiver, one or more amplifiers, a tuner, one or more oscillators, a digital signal processor, a CODEC chipset, a subscriber identity module (SIM) card, memory, and so forth. RF circuitry 208 optionally communicates with networks, such as the Internet, also referred to as the World Wide Web (WWW), an intranet and/or a wireless network, such as a cellular telephone network, a wireless local area network (LAN) and/or a metropolitan area network (MAN), and other devices by wireless communication. The RF circuitry 208 optionally includes well-known circuitry for detecting near field communication (NFC) fields, such as by a short-range communication radio. The wireless communication optionally uses any of a plurality of communications standards, protocols, and technologies, including but not limited to Global System for Mobile Communications (GSM), Enhanced Data GSM Environment (EDGE), high-speed downlink packet access (HSDPA), high-speed uplink packet access (HSUPA), Evolution, Data-Only (EV-DO), HSPA, HSPA+, Dual-Cell HSPA (DC-HSPDA), long term evolution (LTE), near field communication (NFC), wideband code division multiple access (W-CDMA), code division multiple access (CDMA), time division multiple access (TDMA), Bluetooth, Blu-

tooth Low Energy (BTLE), Wireless Fidelity (Wi-Fi) (e.g., IEEE 802.11a, IEEE 802.11b, IEEE 802.11g, IEEE 802.11n, and/or IEEE 802.11ac), voice over Internet Protocol (VoIP), Wi-MAX, a protocol for e mail (e.g., Internet message access protocol (IMAP) and/or post office protocol (POP)), instant messaging (e.g., extensible messaging and presence protocol (XMPP), Session Initiation Protocol for Instant Messaging and Presence Leveraging Extensions (SIMPLE), Instant Messaging and Presence Service (IMPS)), and/or Short Message Service (SMS), or any other suitable communication protocol, including communication protocols not yet developed as of the filing date of this document.

[0060] Audio circuitry 210, speaker 211, and microphone 213 provide an audio interface between a user and device 200. Audio circuitry 210 receives audio data from peripherals interface 218, converts the audio data to an electrical signal, and transmits the electrical signal to speaker 211. Speaker 211 converts the electrical signal to human-audible sound waves. Audio circuitry 210 also receives electrical signals converted by microphone 213 from sound waves. Audio circuitry 210 converts the electrical signal to audio data and transmits the audio data to peripherals interface 218 for processing. Audio data may be retrieved from and/or transmitted to memory 202 and/or RF circuitry 208 by peripherals interface 218. In some embodiments, audio circuitry 210 also includes a headset jack (e.g., 312, FIG. 3). The headset jack provides an interface between audio circuitry 210 and removable audio input/output peripherals, such as output-only headphones or a headset with both output (e.g., a headphone for one or both ears) and input (e.g., a microphone).

[0061] I/O subsystem 206 couples input/output peripherals on device 200, such as touch screen 212 and other input control devices 216, to peripherals interface 218. I/O subsystem 206 optionally includes display controller 256, optical sensor controller 258, intensity sensor controller 259, haptic feedback controller 261, and one or more input controllers 260 for other input or control devices. The one or more input controllers 260 receive/send electrical signals from/to other input control devices 216. The other input control devices 216 optionally include physical buttons (e.g., push buttons, rocker buttons, etc.), dials, slider switches, joysticks, click wheels, and so forth. In some alternate embodiments, input controller(s) 260 are, optionally, coupled to any (or none) of the following: a keyboard, an infrared port, a USB port, and a pointer device such as a mouse. The one or more buttons (e.g., 308, FIG. 3) optionally include an up/down button for volume control of speaker 211 and/or microphone 213. The one or more buttons optionally include a push button (e.g., 306, FIG. 3).

[0062] A quick press of the push button may disengage a lock of touch screen 212 or begin a process that uses gestures on the touch screen to unlock the device, as described in U.S. patent application Ser. No. 11/322,549, “Unlocking a Device by Performing Gestures on an Unlock Image,” filed Dec. 23, 2005, U.S. Pat. No. 7,657,849, which is hereby incorporated by reference in its entirety. A longer press of the push button (e.g., 306) may turn power to device 200 on or off. The user may be able to customize a functionality of one or more of the buttons. Touch screen 212 is used to implement virtual or soft buttons and one or more soft keyboards.

[0063] Touch-sensitive display 212 provides an input interface and an output interface between the device and a

user. Display controller 256 receives and/or sends electrical signals from/to touch screen 212. Touch screen 212 displays visual output to the user. The visual output may include graphics, text, icons, video, and any combination thereof (collectively termed "graphics"). In some embodiments, some or all of the visual output may correspond to user-interface objects.

[0064] Touch screen 212 has a touch-sensitive surface, sensor, or set of sensors that accepts input from the user based on haptic and/or tactile contact. Touch screen 212 and display controller 256 (along with any associated modules and/or sets of instructions in memory 202) detect contact (and any movement or breaking of the contact) on touch screen 212 and convert the detected contact into interaction with user-interface objects (e.g., one or more soft keys, icons, web pages, or images) that are displayed on touch screen 212. In an exemplary embodiment, a point of contact between touch screen 212 and the user corresponds to a finger of the user.

[0065] Touch screen 212 may use LCD (liquid crystal display) technology, LPD (light emitting polymer display) technology, or LED (light emitting diode) technology, although other display technologies may be used in other embodiments. Touch screen 212 and display controller 256 may detect contact and any movement or breaking thereof using any of a plurality of touch sensing technologies now known or later developed, including but not limited to capacitive, resistive, infrared, and surface acoustic wave technologies, as well as other proximity sensor arrays or other elements for determining one or more points of contact with touch screen 212. In an exemplary embodiment, projected mutual capacitance sensing technology is used, such as that found in the iPhone® and iPod Touch® from Apple Inc. of Cupertino, California.

[0066] A touch-sensitive display in some embodiments of touch screen 212 may be analogous to the multi-touch sensitive touchpads described in the following U.S. Pat. No. 6,323,846 (Westerman et al.), U.S. Pat. No. 6,570,557 (Westerman et al.), and/or U.S. Pat. No. 6,677,932 (Westerman), and/or U.S. Patent Publication 2002/0015024A1, each of which is hereby incorporated by reference in its entirety. However, touch screen 212 displays visual output from device 200, whereas touch-sensitive touchpads do not provide visual output.

[0067] A touch-sensitive display in some embodiments of touch screen 212 may be as described in the following applications: (1) U.S. patent application Ser. No. 11/381,313, "Multipoint Touch Surface Controller," filed May 2, 2006; (2) U.S. patent application Ser. No. 10/840,862, "Multipoint Touchscreen," filed May 6, 2004; (3) U.S. patent application Ser. No. 10/903,964, "Gestures For Touch Sensitive Input Devices," filed Jul. 30, 2004; (4) U.S. patent application Ser. No. 11/048,264, "Gestures For Touch Sensitive Input Devices," filed Jan. 31, 2005; (5) U.S. patent application Ser. No. 11/038,590, "Mode-Based Graphical User Interfaces For Touch Sensitive Input Devices," filed Jan. 18, 2005; (6) U.S. patent application Ser. No. 11/228,758, "Virtual Input Device Placement On A Touch Screen User Interface," filed Sep. 16, 2005; (7) U.S. patent application Ser. No. 11/228,700, "Operation Of A Computer With A Touch Screen Interface," filed Sep. 16, 2005; (8) U.S. patent application Ser. No. 11/228,737, "Activating Virtual Keys Of A Touch-Screen Virtual Keyboard," filed Sep. 16, 2005; and (9) U.S. patent application Ser. No. 11/367,749,

"Multi-Functional Hand-Held Device," filed Mar. 3, 2006. All of these applications are incorporated by reference herein in their entirety.

[0068] Touch screen 212 may have a video resolution in excess of 100 dpi. In some embodiments, the touch screen has a video resolution of approximately 160 dpi. The user may make contact with touch screen 212 using any suitable object or appendage, such as a stylus, a finger, and so forth. In some embodiments, the user interface is designed to work primarily with finger-based contacts and gestures, which can be less precise than stylus-based input due to the larger area of contact of a finger on the touch screen. In some embodiments, the device translates the rough finger-based input into a precise pointer/cursor position or command for performing the actions desired by the user.

[0069] In some embodiments, in addition to the touch screen, device 200 may include a touchpad (not shown) for activating or deactivating particular functions. In some embodiments, the touchpad is a touch-sensitive area of the device that, unlike the touch screen, does not display visual output. The touchpad may be a touch-sensitive surface that is separate from touch screen 212 or an extension of the touch-sensitive surface formed by the touch screen.

[0070] Device 200 also includes power system 262 for powering the various components. Power system 262 may include a power management system, one or more power sources (e.g., battery, alternating current (AC)), a recharging system, a power failure detection circuit, a power converter or inverter, a power status indicator (e.g., a light-emitting diode (LED)) and any other components associated with the generation, management and distribution of power in portable devices.

[0071] Device 200 may also include one or more optical sensors 264. FIG. 2A shows an optical sensor coupled to optical sensor controller 258 in I/O subsystem 206. Optical sensor 264 may include charge-coupled device (CCD) or complementary metal-oxide semiconductor (CMOS) phototransistors. Optical sensor 264 receives light from the environment, projected through one or more lenses, and converts the light to data representing an image. In conjunction with imaging module 243 (also called a camera module), optical sensor 264 may capture still images or video. In some embodiments, an optical sensor is located on the back of device 200, opposite touch screen display 212 on the front of the device so that the touch screen display may be used as a viewfinder for still and/or video image acquisition. In some embodiments, an optical sensor is located on the front of the device so that the user's image may be obtained for video conferencing while the user views the other video conference participants on the touch screen display. In some embodiments, the position of optical sensor 264 can be changed by the user (e.g., by rotating the lens and the sensor in the device housing) so that a single optical sensor 264 may be used along with the touch screen display for both video conferencing and still and/or video image acquisition.

[0072] Device 200 optionally also includes one or more contact intensity sensors 265. FIG. 2A shows a contact intensity sensor coupled to intensity sensor controller 259 in I/O subsystem 206. Contact intensity sensor 265 optionally includes one or more piezoresistive strain gauges, capacitive force sensors, electric force sensors, piezoelectric force sensors, optical force sensors, capacitive touch-sensitive surfaces, or other intensity sensors (e.g., sensors used to measure the force (or pressure) of a contact on a touch-

sensitive surface). Contact intensity sensor 265 receives contact intensity information (e.g., pressure information or a proxy for pressure information) from the environment. In some embodiments, at least one contact intensity sensor is collocated with, or proximate to, a touch-sensitive surface (e.g., touch-sensitive display system 212). In some embodiments, at least one contact intensity sensor is located on the back of device 200, opposite touch screen display 212, which is located on the front of device 200.

[0073] Device 200 may also include one or more proximity sensors 266. FIG. 2A shows proximity sensor 266 coupled to peripherals interface 218. Alternately, proximity sensor 266 may be coupled to input controller 260 in I/O subsystem 206. Proximity sensor 266 may perform as described in U.S. patent application Ser. No. 11/241,839, "Proximity Detector In Handheld Device"; Ser. No. 11/240,788, "Proximity Detector In Handheld Device"; Ser. No. 11/620,702, "Using Ambient Light Sensor To Augment Proximity Sensor Output"; Ser. No. 11/586,862, "Automated Response To And Sensing Of User Activity In Portable Devices"; and Ser. No. 11/638,251, "Methods And Systems For Automatic Configuration Of Peripherals," which are hereby incorporated by reference in their entirety. In some embodiments, the proximity sensor turns off and disables touch screen 212 when the multifunction device is placed near the user's ear (e.g., when the user is making a phone call).

[0074] Device 200 optionally also includes one or more tactile output generators 267. FIG. 2A shows a tactile output generator coupled to haptic feedback controller 261 in I/O subsystem 206. Tactile output generator 267 optionally includes one or more electroacoustic devices such as speakers or other audio components and/or electromechanical devices that convert energy into linear motion such as a motor, solenoid, electroactive polymer, piezoelectric actuator, electrostatic actuator, or other tactile output generating component (e.g., a component that converts electrical signals into tactile outputs on the device). Contact intensity sensor 265 receives tactile feedback generation instructions from haptic feedback module 233 and generates tactile outputs on device 200 that are capable of being sensed by a user of device 200. In some embodiments, at least one tactile output generator is collocated with, or proximate to, a touch-sensitive surface (e.g., touch-sensitive display system 212) and, optionally, generates a tactile output by moving the touch-sensitive surface vertically (e.g., in/out of a surface of device 200) or laterally (e.g., back and forth in the same plane as a surface of device 200). In some embodiments, at least one tactile output generator sensor is located on the back of device 200, opposite touch screen display 212, which is located on the front of device 200.

[0075] Device 200 may also include one or more accelerometers 268. FIG. 2A shows accelerometer 268 coupled to peripherals interface 218. Alternately, accelerometer 268 may be coupled to an input controller 260 in I/O subsystem 206. Accelerometer 268 may perform as described in U.S. Patent Publication No. 20050190059, "Acceleration-based Theft Detection System for Portable Electronic Devices," and U.S. Patent Publication No. 20060017692, "Methods And Apparatuses For Operating A Portable Device Based On An Accelerometer," both of which are incorporated by reference herein in their entirety. In some embodiments, information is displayed on the touch screen display in a portrait view or a landscape view based on an analysis of

data received from the one or more accelerometers. Device 200 optionally includes, in addition to accelerometer(s) 268, a magnetometer (not shown) and a GPS (or GLONASS or other global navigation system) receiver (not shown) for obtaining information concerning the location and orientation (e.g., portrait or landscape) of device 200.

[0076] In some embodiments, the software components stored in memory 202 include operating system 226, communication module (or set of instructions) 228, contact/motion module (or set of instructions) 230, graphics module (or set of instructions) 232, text input module (or set of instructions) 234, Global Positioning System (GPS) module (or set of instructions) 235, Digital Assistant Client Module 229, and applications (or sets of instructions) 236. Further, memory 202 can store data and models, such as user data and models 231. Furthermore, in some embodiments, memory 202 (FIG. 2A) or 470 (FIG. 4) stores device/global internal state 257, as shown in FIGS. 2A and 4. Device/global internal state 257 includes one or more of: active application state, indicating which applications, if any, are currently active; display state, indicating what applications, views or other information occupy various regions of touch screen display 212; sensor state, including information obtained from the device's various sensors and input control devices 216; and location information concerning the device's location and/or attitude.

[0077] Operating system 226 (e.g., Darwin, RTXC, LINUX, UNIX, OS X, iOS, WINDOWS, or an embedded operating system such as VxWorks) includes various software components and/or drivers for controlling and managing general system tasks (e.g., memory management, storage device control, power management, etc.) and facilitates communication between various hardware and software components.

[0078] Communication module 228 facilitates communication with other devices over one or more external ports 224 and also includes various software components for handling data received by RF circuitry 208 and/or external port 224. External port 224 (e.g., Universal Serial Bus (USB), FIREWIRE, etc.) is adapted for coupling directly to other devices or indirectly over a network (e.g., the Internet, wireless LAN, etc.). In some embodiments, the external port is a multi-pin (e.g., 30-pin) connector that is the same as, or similar to and/or compatible with, the 30-pin connector used on iPod® (trademark of Apple Inc.) devices.

[0079] Contact/motion module 230 optionally detects contact with touch screen 212 (in conjunction with display controller 256) and other touch-sensitive devices (e.g., a touchpad or physical click wheel). Contact/motion module 230 includes various software components for performing various operations related to detection of contact, such as determining if contact has occurred (e.g., detecting a finger-down event), determining an intensity of the contact (e.g., the force or pressure of the contact or a substitute for the force or pressure of the contact), determining if there is movement of the contact and tracking the movement across the touch-sensitive surface (e.g., detecting one or more finger-dragging events), and determining if the contact has ceased (e.g., detecting a finger-up event or a break in contact). Contact/motion module 230 receives contact data from the touch-sensitive surface. Determining movement of the point of contact, which is represented by a series of contact data, optionally includes determining speed (magnitude), velocity (magnitude and direction), and/or an accel-

eration (a change in magnitude and/or direction) of the point of contact. These operations are, optionally, applied to single contacts (e.g., one finger contacts) or to multiple simultaneous contacts (e.g., “multitouch”/multiple finger contacts). In some embodiments, contact/motion module 230 and display controller 256 detect contact on a touchpad.

[0080] In some embodiments, contact/motion module 230 uses a set of one or more intensity thresholds to determine whether an operation has been performed by a user (e.g., to determine whether a user has “clicked” on an icon). In some embodiments, at least a subset of the intensity thresholds are determined in accordance with software parameters (e.g., the intensity thresholds are not determined by the activation thresholds of particular physical actuators and can be adjusted without changing the physical hardware of device 200). For example, a mouse “click” threshold of a trackpad or touch screen display can be set to any of a large range of predefined threshold values without changing the trackpad or touch screen display hardware. Additionally, in some implementations, a user of the device is provided with software settings for adjusting one or more of the set of intensity thresholds (e.g., by adjusting individual intensity thresholds and/or by adjusting a plurality of intensity thresholds at once with a system-level click “intensity” parameter).

[0081] Contact/motion module 230 optionally detects a gesture input by a user. Different gestures on the touch-sensitive surface have different contact patterns (e.g., different motions, timings, and/or intensities of detected contacts). Thus, a gesture is, optionally, detected by detecting a particular contact pattern. For example, detecting a finger tap gesture includes detecting a finger-down event followed by detecting a finger-up (liftoff) event at the same position (or substantially the same position) as the finger-down event (e.g., at the position of an icon). As another example, detecting a finger swipe gesture on the touch-sensitive surface includes detecting a finger-down event followed by detecting one or more finger-dragging events, and subsequently followed by detecting a finger-up (liftoff) event.

[0082] Graphics module 232 includes various known software components for rendering and displaying graphics on touch screen 212 or other display, including components for changing the visual impact (e.g., brightness, transparency, saturation, contrast, or other visual property) of graphics that are displayed. As used herein, the term “graphics” includes any object that can be displayed to a user, including, without limitation, text, web pages, icons (such as user-interface objects including soft keys), digital images, videos, animations, and the like.

[0083] In some embodiments, graphics module 232 stores data representing graphics to be used. Each graphic is, optionally, assigned a corresponding code. Graphics module 232 receives, from applications etc., one or more codes specifying graphics to be displayed along with, if necessary, coordinate data and other graphic property data, and then generates screen image data to output to display controller 256.

[0084] Haptic feedback module 233 includes various software components for generating instructions used by tactile output generator(s) 267 to produce tactile outputs at one or more locations on device 200 in response to user interactions with device 200.

[0085] Text input module 234, which may be a component of graphics module 232, provides soft keyboards for entering text in various applications (e.g., contacts module 237,

e-mail client module 240, IM module 241, browser module 247, and any other application that needs text input).

[0086] GPS module 235 determines the location of the device and provides this information for use in various applications (e.g., to telephone module 238 for use in location-based dialing; to camera module 243 as picture/video metadata; and to applications that provide location-based services such as weather widgets, local yellow page widgets, and map/navigation widgets).

[0087] Digital assistant client module 229 can include various client-side digital assistant instructions to provide the client-side functionalities of the digital assistant. For example, digital assistant client module 229 can be capable of accepting voice input (e.g., speech input), text input, touch input, and/or gestural input through various user interfaces (e.g., microphone 213, accelerometer(s) 268, touch-sensitive display system 212, optical sensor(s) 229, other input control devices 216, etc.) of portable multifunction device 200. Digital assistant client module 229 can also be capable of providing output in audio (e.g., speech output), visual, and/or tactile forms through various output interfaces (e.g., speaker 211, touch-sensitive display system 212, tactile output generator(s) 267, etc.) of portable multifunction device 200. For example, output can be provided as voice, sound, alerts, text messages, menus, graphics, videos, animations, vibrations, and/or combinations of two or more of the above. During operation, digital assistant client module 229 can communicate with DA server 106 using RF circuitry 208.

[0088] User data and models 231 can include various data associated with the user (e.g., user-specific vocabulary data, user preference data, user-specified name pronunciations, data from the user’s electronic address book, to-do lists, shopping lists, etc.) to provide the client-side functionalities of the digital assistant. Further, user data and models 231 can include various models (e.g., speech recognition models, statistical language models, natural language processing models, ontology, task flow models, service models, etc.) for processing user input and determining user intent.

[0089] In some examples, digital assistant client module 229 can utilize the various sensors, subsystems, and peripheral devices of portable multifunction device 200 to gather additional information from the surrounding environment of the portable multifunction device 200 to establish a context associated with a user, the current user interaction, and/or the current user input. In some examples, digital assistant client module 229 can provide the contextual information or a subset thereof with the user input to DA server 106 to help infer the user’s intent. In some examples, the digital assistant can also use the contextual information to determine how to prepare and deliver outputs to the user. Contextual information can be referred to as context data.

[0090] In some examples, the contextual information that accompanies the user input can include sensor information, e.g., lighting, ambient noise, ambient temperature, images or videos of the surrounding environment, etc. In some examples, the contextual information can also include the physical state of the device, e.g., device orientation, device location, device temperature, power level, speed, acceleration, motion patterns, cellular signals strength, etc. In some examples, information related to the software state of DA server 106, e.g., running processes, installed programs, past and present network activities, background services, error logs, resources usage, etc., and of portable multifunction

device 200 can be provided to DA server 106 as contextual information associated with a user input.

[0091] In some examples, the digital assistant client module 229 can selectively provide information (e.g., user data 231) stored on the portable multifunction device 200 in response to requests from DA server 106. In some examples, digital assistant client module 229 can also elicit additional input from the user via a natural language dialogue or other user interfaces upon request by DA server 106. Digital assistant client module 229 can pass the additional input to DA server 106 to help DA server 106 in intent deduction and/or fulfillment of the user's intent expressed in the user request.

[0092] A more detailed description of a digital assistant is described below with reference to FIGS. 7A-7C. It should be recognized that digital assistant client module 229 can include any number of the sub-modules of digital assistant module 726 described below.

[0093] Applications 236 may include the following modules (or sets of instructions), or a subset or superset thereof:

- [0094] Contacts module 237 (sometimes called an address book or contact list);
- [0095] Telephone module 238;
- [0096] Video conference module 239;
- [0097] E-mail client module 240;
- [0098] Instant messaging (IM) module 241;
- [0099] Workout support module 242;
- [0100] Camera module 243 for still and/or video images;
- [0101] Image management module 244;
- [0102] Video player module;
- [0103] Music player module;
- [0104] Browser module 247;
- [0105] Calendar module 248;
- [0106] Widget modules 249, which may include one or more of: weather widget 249-1, stocks widget 249-2, calculator widget 249-3, alarm clock widget 249-4, dictionary widget 249-5, and other widgets obtained by the user, as well as user-created widgets 249-6;
- [0107] Widget creator module 250 for making user-created widgets 249-6;
- [0108] Search module 251;
- [0109] Video and music player module 252, which merges video player module and music player module;
- [0110] Notes module 253;
- [0111] Map module 254; and/or
- [0112] Online video module 255.

[0113] Examples of other applications 236 that may be stored in memory 202 include other word processing applications, other image editing applications, drawing applications, presentation applications, JAVA-enabled applications, encryption, digital rights management, voice recognition, and voice replication.

[0114] In conjunction with touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, contacts module 237 may be used to manage an address book or contact list (e.g., stored in application internal state 292 of contacts module 237 in memory 202 or memory 470), including: adding name(s) to the address book; deleting name(s) from the address book; associating telephone number(s), email address(es), physical address(es) or other information with a name; associating an image with a name; categorizing and sorting names; providing telephone numbers or email addresses to initiate

and/or facilitate communications by telephone module 238, video conference module 239, e-mail client module 240, or IM module 241; and so forth.

[0115] In conjunction with RF circuitry 208, audio circuitry 210, speaker 211, microphone 213, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, telephone module 238 may be used to enter a sequence of characters corresponding to a telephone number, access one or more telephone numbers in contacts module 237, modify a telephone number that has been entered, dial a respective telephone number, conduct a conversation, and disconnect or hang up when the conversation is completed. As noted above, the wireless communication may use any of a plurality of communications standards, protocols, and technologies.

[0116] In conjunction with RF circuitry 208, audio circuitry 210, speaker 211, microphone 213, touch screen 212, display controller 256, optical sensor 264, optical sensor controller 258, contact/motion module 230, graphics module 232, text input module 234, contacts module 237, and telephone module 238, video conference module 239 includes executable instructions to initiate, conduct, and terminate a video conference between a user and one or more other participants in accordance with user instructions.

[0117] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, email client module 240 includes executable instructions to create, send, receive, and manage email in response to user instructions. In conjunction with image management module 244, email client module 240 makes it very easy to create and send emails with still or video images taken with camera module 243.

[0118] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, the instant messaging module 241 includes executable instructions to enter a sequence of characters corresponding to an instant message, to modify previously entered characters, to transmit a respective instant message (for example, using a Short Message Service (SMS) or Multimedia Message Service (MMS) protocol for telephony-based instant messages or using XMPP, SIMPLE, or IMPS for Internet-based instant messages), to receive instant messages, and to view received instant messages. In some embodiments, transmitted and/or received instant messages may include graphics, photos, audio files, video files and/or other attachments as are supported in an MMS and/or an Enhanced Messaging Service (EMS). As used herein, "instant messaging" refers to both telephony-based messages (e.g., messages sent using SMS or MMS) and Internet-based messages (e.g., messages sent using XMPP, SIMPLE, or IMPS).

[0119] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, text input module 234, GPS module 235, map module 254, and music player module, workout support module 242 includes executable instructions to create workouts (e.g., with time, distance, and/or calorie burning goals); communicate with workout sensors (sports devices); receive workout sensor data; calibrate sensors used to monitor a workout; select and play music for a workout; and display, store, and transmit workout data.

[0120] In conjunction with touch screen 212, display controller 256, optical sensor(s) 264, optical sensor control-

ler 258, contact/motion module 230, graphics module 232, and image management module 244, camera module 243 includes executable instructions to capture still images or video (including a video stream) and store them into memory 202, modify characteristics of a still image or video, or delete a still image or video from memory 202.

[0121] In conjunction with touch screen 212, display controller 256, contact/motion module 230, graphics module 232, text input module 234, and camera module 243, image management module 244 includes executable instructions to arrange, modify (e.g., edit), or otherwise manipulate, label, delete, present (e.g., in a digital slide show or album), and store still and/or video images.

[0122] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, browser module 247 includes executable instructions to browse the Internet in accordance with user instructions, including searching, linking to, receiving, and displaying web pages or portions thereof, as well as attachments and other files linked to web pages.

[0123] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, text input module 234, email client module 240, and browser module 247, calendar module 248 includes executable instructions to create, display, modify, and store calendars and data associated with calendars (e.g., calendar entries, to-do lists, etc.) in accordance with user instructions.

[0124] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, text input module 234, and browser module 247, widget modules 249 are mini-applications that may be downloaded and used by a user (e.g., weather widget 249-1, stocks widget 249-2, calculator widget 249-3, alarm clock widget 249-4, and dictionary widget 249-5) or created by the user (e.g., user-created widget 249-6). In some embodiments, a widget includes an HTML (Hypertext Markup Language) file, a CSS (Cascading Style Sheets) file, and a JavaScript file. In some embodiments, a widget includes an XML (Extensible Markup Language) file and a JavaScript file (e.g., Yahoo!Widgets).

[0125] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, text input module 234, and browser module 247, the widget creator module 250 may be used by a user to create widgets (e.g., turning a user-specified portion of a web page into a widget).

[0126] In conjunction with touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, search module 251 includes executable instructions to search for text, music, sound, image, video, and/or other files in memory 202 that match one or more search criteria (e.g., one or more user-specified search terms) in accordance with user instructions.

[0127] In conjunction with touch screen 212, display controller 256, contact/motion module 230, graphics module 232, audio circuitry 210, speaker 211, RF circuitry 208, and browser module 247, video and music player module 252 includes executable instructions that allow the user to download and play back recorded music and other sound files stored in one or more file formats, such as MP3 or AAC files, and executable instructions to display, present, or otherwise play back videos (e.g., on touch screen 212 or on an external,

connected display via external port 224). In some embodiments, device 200 optionally includes the functionality of an MP3 player, such as an iPod (trademark of Apple Inc.).

[0128] In conjunction with touch screen 212, display controller 256, contact/motion module 230, graphics module 232, and text input module 234, notes module 253 includes executable instructions to create and manage notes, to-do lists, and the like in accordance with user instructions.

[0129] In conjunction with RF circuitry 208, touch screen 212, display controller 256, contact/motion module 230, graphics module 232, text input module 234, GPS module 235, and browser module 247, map module 254 may be used to receive, display, modify, and store maps and data associated with maps (e.g., driving directions, data on stores and other points of interest at or near a particular location, and other location-based data) in accordance with user instructions.

[0130] In conjunction with touch screen 212, display controller 256, contact/motion module 230, graphics module 232, audio circuitry 210, speaker 211, RF circuitry 208, text input module 234, email client module 240, and browser module 247, online video module 255 includes instructions that allow the user to access, browse, receive (e.g., by streaming and/or download), play back (e.g., on the touch screen or on an external, connected display via external port 224), send an email with a link to a particular online video, and otherwise manage online videos in one or more file formats, such as H.264. In some embodiments, instant messaging module 241, rather than email client module 240, is used to send a link to a particular online video. Additional description of the online video application can be found in U.S. Provisional Patent Application No. 60/936,562, "Portable Multifunction Device, Method, and Graphical User Interface for Playing Online Videos," filed Jun. 20, 2007, and U.S. patent application Ser. No. 11/968,067, "Portable Multifunction Device, Method, and Graphical User Interface for Playing Online Videos," filed Dec. 31, 2007, the contents of which are hereby incorporated by reference in their entirety.

[0131] Each of the above-identified modules and applications corresponds to a set of executable instructions for performing one or more functions described above and the methods described in this application (e.g., the computer-implemented methods and other information processing methods described herein). These modules (e.g., sets of instructions) need not be implemented as separate software programs, procedures, or modules, and thus various subsets of these modules may be combined or otherwise rearranged in various embodiments. For example, video player module may be combined with music player module into a single module (e.g., video and music player module 252, FIG. 2A). In some embodiments, memory 202 may store a subset of the modules and data structures identified above. Furthermore, memory 202 may store additional modules and data structures not described above.

[0132] In some embodiments, device 200 is a device where operation of a predefined set of functions on the device is performed exclusively through a touch screen and/or a touchpad. By using a touch screen and/or a touchpad as the primary input control device for operation of device 200, the number of physical input control devices (such as push buttons, dials, and the like) on device 200 may be reduced.

[0133] The predefined set of functions that are performed exclusively through a touch screen and/or a touchpad optionally include navigation between user interfaces. In some embodiments, the touchpad, when touched by the user, navigates device 200 to a main, home, or root menu from any user interface that is displayed on device 200. In such embodiments, a “menu button” is implemented using a touchpad. In some other embodiments, the menu button is a physical push button or other physical input control device instead of a touchpad.

[0134] FIG. 2B is a block diagram illustrating exemplary components for event handling in accordance with some embodiments. In some embodiments, memory 202 (FIG. 2A) or 470 (FIG. 4) includes event sorter 270 (e.g., in operating system 226) and a respective application 236-1 (e.g., any of the aforementioned applications 237-251, 255, 480-490).

[0135] Event sorter 270 receives event information and determines the application 236-1 and application view 291 of application 236-1 to which to deliver the event information. Event sorter 270 includes event monitor 271 and event dispatcher module 274. In some embodiments, application 236-1 includes application internal state 292, which indicates the current application view(s) displayed on touch-sensitive display 212 when the application is active or executing. In some embodiments, device/global internal state 257 is used by event sorter 270 to determine which application(s) is (are) currently active, and application internal state 292 is used by event sorter 270 to determine application views 291 to which to deliver event information.

[0136] In some embodiments, application internal state 292 includes additional information, such as one or more of: resume information to be used when application 236-1 resumes execution, user interface state information that indicates information being displayed or that is ready for display by application 236-1, a state queue for enabling the user to go back to a prior state or view of application 236-1, and a redo/undo queue of previous actions taken by the user.

[0137] Event monitor 271 receives event information from peripherals interface 218. Event information includes information about a sub-event (e.g., a user touch on touch-sensitive display 212, as part of a multi-touch gesture). Peripherals interface 218 transmits information it receives from I/O subsystem 206 or a sensor, such as proximity sensor 266, accelerometer(s) 268, and/or microphone 213 (through audio circuitry 210). Information that peripherals interface 218 receives from I/O subsystem 206 includes information from touch-sensitive display 212 or a touch-sensitive surface.

[0138] In some embodiments, event monitor 271 sends requests to the peripherals interface 218 at predetermined intervals. In response, peripherals interface 218 transmits event information. In other embodiments, peripherals interface 218 transmits event information only when there is a significant event (e.g., receiving an input above a predetermined noise threshold and/or for more than a predetermined duration).

[0139] In some embodiments, event sorter 270 also includes a hit view determination module 272 and/or an active event recognizer determination module 273.

[0140] Hit view determination module 272 provides software procedures for determining where a sub-event has taken place within one or more views when touch-sensitive

display 212 displays more than one view. Views are made up of controls and other elements that a user can see on the display.

[0141] Another aspect of the user interface associated with an application is a set of views, sometimes herein called application views or user interface windows, in which information is displayed and touch-based gestures occur. The application views (of a respective application) in which a touch is detected may correspond to programmatic levels within a programmatic or view hierarchy of the application. For example, the lowest level view in which a touch is detected may be called the hit view, and the set of events that are recognized as proper inputs may be determined based, at least in part, on the hit view of the initial touch that begins a touch-based gesture.

[0142] Hit view determination module 272 receives information related to sub events of a touch-based gesture. When an application has multiple views organized in a hierarchy, hit view determination module 272 identifies a hit view as the lowest view in the hierarchy which should handle the sub-event. In most circumstances, the hit view is the lowest level view in which an initiating sub-event occurs (e.g., the first sub-event in the sequence of sub-events that form an event or potential event). Once the hit view is identified by the hit view determination module 272, the hit view typically receives all sub-events related to the same touch or input source for which it was identified as the hit view.

[0143] Active event recognizer determination module 273 determines which view or views within a view hierarchy should receive a particular sequence of sub-events. In some embodiments, active event recognizer determination module 273 determines that only the hit view should receive a particular sequence of sub-events. In other embodiments, active event recognizer determination module 273 determines that all views that include the physical location of a sub-event are actively involved views, and therefore determines that all actively involved views should receive a particular sequence of sub-events. In other embodiments, even if touch sub-events were entirely confined to the area associated with one particular view, views higher in the hierarchy would still remain as actively involved views.

[0144] Event dispatcher module 274 dispatches the event information to an event recognizer (e.g., event recognizer 280). In embodiments including active event recognizer determination module 273, event dispatcher module 274 delivers the event information to an event recognizer determined by active event recognizer determination module 273. In some embodiments, event dispatcher module 274 stores in an event queue the event information, which is retrieved by a respective event receiver 282.

[0145] In some embodiments, operating system 226 includes event sorter 270. Alternatively, application 236-1 includes event sorter 270. In yet other embodiments, event sorter 270 is a stand-alone module, or a part of another module stored in memory 202, such as contact/motion module 230.

[0146] In some embodiments, application 236-1 includes a plurality of event handlers 290 and one or more application views 291, each of which includes instructions for handling touch events that occur within a respective view of the application's user interface. Each application view 291 of the application 236-1 includes one or more event recognizers 280. Typically, a respective application view 291 includes a plurality of event recognizers 280. In other

embodiments, one or more of event recognizers 280 are part of a separate module, such as a user interface kit (not shown) or a higher level object from which application 236-1 inherits methods and other properties. In some embodiments, a respective event handler 290 includes one or more of: data updater 276, object updater 277, GUI updater 278, and/or event data 279 received from event sorter 270. Event handler 290 may utilize or call data updater 276, object updater 277, or GUI updater 278 to update the application internal state 292. Alternatively, one or more of the application views 291 include one or more respective event handlers 290. Also, in some embodiments, one or more of data updater 276, object updater 277, and GUI updater 278 are included in a respective application view 291.

[0147] A respective event recognizer 280 receives event information (e.g., event data 279) from event sorter 270 and identifies an event from the event information. Event recognizer 280 includes event receiver 282 and event comparator 284. In some embodiments, event recognizer 280 also includes at least a subset of: metadata 283, and event delivery instructions 288 (which may include sub-event delivery instructions).

[0148] Event receiver 282 receives event information from event sorter 270. The event information includes information about a sub-event, for example, a touch or a touch movement. Depending on the sub-event, the event information also includes additional information, such as location of the sub-event. When the sub-event concerns motion of a touch, the event information may also include speed and direction of the sub-event. In some embodiments, events include rotation of the device from one orientation to another (e.g., from a portrait orientation to a landscape orientation, or vice versa), and the event information includes corresponding information about the current orientation (also called device attitude) of the device.

[0149] Event comparator 284 compares the event information to predefined event or sub-event definitions and, based on the comparison, determines an event or sub event, or determines or updates the state of an event or sub-event. In some embodiments, event comparator 284 includes event definitions 286. Event definitions 286 contain definitions of events (e.g., predefined sequences of sub-events), for example, event 1 (287-1), event 2 (287-2), and others. In some embodiments, sub-events in an event (287) include, for example, touch begin, touch end, touch movement, touch cancellation, and multiple touching. In one example, the definition for event 1 (287-1) is a double tap on a displayed object. The double tap, for example, comprises a first touch (touch begin) on the displayed object for a predetermined phase, a first liftoff (touch end) for a predetermined phase, a second touch (touch begin) on the displayed object for a predetermined phase, and a second liftoff (touch end) for a predetermined phase. In another example, the definition for event 2 (287-2) is a dragging on a displayed object. The dragging, for example, comprises a touch (or contact) on the displayed object for a predetermined phase, a movement of the touch across touch-sensitive display 212, and liftoff of the touch (touch end). In some embodiments, the event also includes information for one or more associated event handlers 290.

[0150] In some embodiments, event definition 287 includes a definition of an event for a respective user-interface object. In some embodiments, event comparator 284 performs a hit test to determine which user-interface

object is associated with a sub-event. For example, in an application view in which three user-interface objects are displayed on touch-sensitive display 212, when a touch is detected on touch-sensitive display 212, event comparator 284 performs a hit test to determine which of the three user-interface objects is associated with the touch (sub-event). If each displayed object is associated with a respective event handler 290, the event comparator uses the result of the hit test to determine which event handler 290 should be activated. For example, event comparator 284 selects an event handler associated with the sub-event and the object triggering the hit test.

[0151] In some embodiments, the definition for a respective event (287) also includes delayed actions that delay delivery of the event information until after it has been determined whether the sequence of sub-events does or does not correspond to the event recognizer's event type.

[0152] When a respective event recognizer 280 determines that the series of sub-events do not match any of the events in event definitions 286, the respective event recognizer 280 enters an event impossible, event failed, or event ended state, after which it disregards subsequent sub-events of the touch-based gesture. In this situation, other event recognizers, if any, that remain active for the hit view continue to track and process sub-events of an ongoing touch-based gesture.

[0153] In some embodiments, a respective event recognizer 280 includes metadata 283 with configurable properties, flags, and/or lists that indicate how the event delivery system should perform sub-event delivery to actively involved event recognizers. In some embodiments, metadata 283 includes configurable properties, flags, and/or lists that indicate how event recognizers may interact, or are enabled to interact, with one another. In some embodiments, metadata 283 includes configurable properties, flags, and/or lists that indicate whether sub-events are delivered to varying levels in the view or programmatic hierarchy.

[0154] In some embodiments, a respective event recognizer 280 activates event handler 290 associated with an event when one or more particular sub-events of an event are recognized. In some embodiments, a respective event recognizer 280 delivers event information associated with the event to event handler 290. Activating an event handler 290 is distinct from sending (and deferred sending) sub-events to a respective hit view. In some embodiments, event recognizer 280 throws a flag associated with the recognized event, and event handler 290 associated with the flag catches the flag and performs a predefined process.

[0155] In some embodiments, event delivery instructions 288 include sub-event delivery instructions that deliver event information about a sub-event without activating an event handler. Instead, the sub-event delivery instructions deliver event information to event handlers associated with the series of sub-events or to actively involved views. Event handlers associated with the series of sub-events or with actively involved views receive the event information and perform a predetermined process.

[0156] In some embodiments, data updater 276 creates and updates data used in application 236-1. For example, data updater 276 updates the telephone number used in contacts module 237, or stores a video file used in video player module. In some embodiments, object updater 277 creates and updates objects used in application 236-1. For example, object updater 277 creates a new user-interface

object or updates the position of a user-interface object. GUI updater 278 updates the GUI. For example, GUI updater 278 prepares display information and sends it to graphics module 232 for display on a touch-sensitive display.

[0157] In some embodiments, event handler(s) 290 includes or has access to data updater 276, object updater 277, and GUI updater 278. In some embodiments, data updater 276, object updater 277, and GUI updater 278 are included in a single module of a respective application 236-1 or application view 291. In other embodiments, they are included in two or more software modules.

[0158] It shall be understood that the foregoing discussion regarding event handling of user touches on touch-sensitive displays also applies to other forms of user inputs to operate multifunction devices 200 with input devices, not all of which are initiated on touch screens. For example, mouse movement and mouse button presses, optionally coordinated with single or multiple keyboard presses or holds; contact movements such as taps, drags, scrolls, etc. on touchpads; pen stylus inputs; movement of the device; oral instructions; detected eye movements; biometric inputs; and/or any combination thereof are optionally utilized as inputs corresponding to sub-events which define an event to be recognized.

[0159] FIG. 3 illustrates a portable multifunction device 200 having a touch screen 212 in accordance with some embodiments. The touch screen optionally displays one or more graphics within user interface (UI) 300. In this embodiment, as well as others described below, a user is enabled to select one or more of the graphics by making a gesture on the graphics, for example, with one or more fingers 302 (not drawn to scale in the figure) or one or more styluses 303 (not drawn to scale in the figure). In some embodiments, selection of one or more graphics occurs when the user breaks contact with the one or more graphics. In some embodiments, the gesture optionally includes one or more taps, one or more swipes (from left to right, right to left, upward and/or downward), and/or a rolling of a finger (from right to left, left to right, upward and/or downward) that has made contact with device 200. In some implementations or circumstances, inadvertent contact with a graphic does not select the graphic. For example, a swipe gesture that sweeps over an application icon optionally does not select the corresponding application when the gesture corresponding to selection is a tap.

[0160] Device 200 may also include one or more physical buttons, such as "home" or menu button 304. As described previously, menu button 304 may be used to navigate to any application 236 in a set of applications that may be executed on device 200. Alternatively, in some embodiments, the menu button is implemented as a soft key in a GUI displayed on touch screen 212.

[0161] In one embodiment, device 200 includes touch screen 212, menu button 304, push button 306 for powering the device on/off and locking the device, volume adjustment button(s) 308, subscriber identity module (SIM) card slot 310, headset jack 312, and docking/charging external port 224. Push button 306 is, optionally, used to turn the power on/off on the device by depressing the button and holding the button in the depressed state for a predefined time interval; to lock the device by depressing the button and releasing the button before the predefined time interval has elapsed; and/or to unlock the device or initiate an unlock process. In an alternative embodiment, device 200 also accepts verbal input for activation or deactivation of some functions

through microphone 213. Device 200 also, optionally, includes one or more contact intensity sensors 265 for detecting intensity of contacts on touch screen 212 and/or one or more tactile output generators 267 for generating tactile outputs for a user of device 200.

[0162] FIG. 4 is a block diagram of an exemplary multifunction device with a display and a touch-sensitive surface in accordance with some embodiments. Device 400 need not be portable. In some embodiments, device 400 is a laptop computer, a desktop computer, a tablet computer, a multimedia player device, a navigation device, an educational device (such as a child's learning toy), a gaming system, or a control device (e.g., a home or industrial controller). Device 400 typically includes one or more processing units (CPUs) 410, one or more network or other communications interfaces 460, memory 470, and one or more communication buses 420 for interconnecting these components. Communication buses 420 optionally include circuitry (sometimes called a chipset) that interconnects and controls communications between system components. Device 400 includes input/output (I/O) interface 430 comprising display 440, which is typically a touch screen display. I/O interface 430 also optionally includes a keyboard and/or mouse (or other pointing device) 450 and touchpad 455, tactile output generator 457 for generating tactile outputs on device 400 (e.g., similar to tactile output generator(s) 267 described above with reference to FIG. 2A), sensors 459 (e.g., optical, acceleration, proximity, touch-sensitive, and/or contact intensity sensors similar to contact intensity sensor(s) 265 described above with reference to FIG. 2A). Memory 470 includes high-speed random access memory, such as DRAM, SRAM, DDR RAM, or other random access solid state memory devices; and optionally includes non-volatile memory, such as one or more magnetic disk storage devices, optical disk storage devices, flash memory devices, or other non-volatile solid state storage devices. Memory 470 optionally includes one or more storage devices remotely located from CPU(s) 410. In some embodiments, memory 470 stores programs, modules, and data structures analogous to the programs, modules, and data structures stored in memory 202 of portable multifunction device 200 (FIG. 2A), or a subset thereof. Furthermore, memory 470 optionally stores additional programs, modules, and data structures not present in memory 202 of portable multifunction device 200. For example, memory 470 of device 400 optionally stores drawing module 480, presentation module 482, word processing module 484, website creation module 486, disk authoring module 488, and/or spreadsheet module 490, while memory 202 of portable multifunction device 200 (FIG. 2A) optionally does not store these modules.

[0163] Each of the above-identified elements in FIG. 4 may be stored in one or more of the previously mentioned memory devices. Each of the above-identified modules corresponds to a set of instructions for performing a function described above. The above-identified modules or programs (e.g., sets of instructions) need not be implemented as separate software programs, procedures, or modules, and thus various subsets of these modules may be combined or otherwise rearranged in various embodiments. In some embodiments, memory 470 may store a subset of the modules and data structures identified above. Furthermore, memory 470 may store additional modules and data structures not described above.

[0164] Attention is now directed towards embodiments of user interfaces that may be implemented on, for example, portable multifunction device 200.

[0165] FIG. 5A illustrates an exemplary user interface for a menu of applications on portable multifunction device 200 in accordance with some embodiments. Similar user interfaces may be implemented on device 400. In some embodiments, user interface 500 includes the following elements, or a subset or superset thereof:

- [0166] Signal strength indicator(s) 502 for wireless communication(s), such as cellular and Wi-Fi signals;
- [0167] Time 504;
- [0168] Bluetooth indicator 505;
- [0169] Battery status indicator 506;
- [0170] Tray 508 with icons for frequently used applications, such as:
 - [0171] Icon 516 for telephone module 238, labeled “Phone,” which optionally includes an indicator 514 of the number of missed calls or voicemail messages;
 - [0172] Icon 518 for email client module 240, labeled “Mail,” which optionally includes an indicator 510 of the number of unread emails;
 - [0173] Icon 520 for browser module 247, labeled “Browser;” and
 - [0174] Icon 522 for video and music player module 252, also referred to as iPod (trademark of Apple Inc.) module 252, labeled “iPod;” and
- [0175] Icons for other applications, such as:
 - [0176] Icon 524 for IM module 241, labeled “Messages;”
 - [0177] Icon 526 for calendar module 248, labeled “Calendar;”
 - [0178] Icon 528 for image management module 244, labeled “Photos;”
 - [0179] Icon 530 for camera module 243, labeled “Camera;”
 - [0180] Icon 532 for online video module 255, labeled “Online Video;”
 - [0181] Icon 534 for stocks widget 249-2, labeled “Stocks;”
 - [0182] Icon 536 for map module 254, labeled “Maps;”
 - [0183] Icon 538 for weather widget 249-1, labeled “Weather;”
 - [0184] Icon 540 for alarm clock widget 249-4, labeled “Clock;”
 - [0185] Icon 542 for workout support module 242, labeled “Workout Support;”
 - [0186] Icon 544 for notes module 253, labeled “Notes;” and
 - [0187] Icon 546 for a settings application or module, labeled “Settings,” which provides access to settings for device 200 and its various applications 236.

[0188] It should be noted that the icon labels illustrated in FIG. 5A are merely exemplary. For example, icon 522 for video and music player module 252 may optionally be labeled “Music” or “Music Player.” Other labels are, optionally, used for various application icons. In some embodiments, a label for a respective application icon includes a name of an application corresponding to the respective application icon. In some embodiments, a label for a particular application icon is distinct from a name of an application corresponding to the particular application icon.

[0189] FIG. 5B illustrates an exemplary user interface on a device (e.g., device 400, FIG. 4) with a touch-sensitive surface 551 (e.g., a tablet or touchpad 455, FIG. 4) that is separate from the display 550 (e.g., touch screen display 212). Device 400 also, optionally, includes one or more contact intensity sensors (e.g., one or more of sensors 457) for detecting intensity of contacts on touch-sensitive surface 551 and/or one or more tactile output generators 459 for generating tactile outputs for a user of device 400.

[0190] Although some of the examples which follow will be given with reference to inputs on touch screen display 212 (where the touch-sensitive surface and the display are combined), in some embodiments, the device detects inputs on a touch-sensitive surface that is separate from the display, as shown in FIG. 5B. In some embodiments, the touch-sensitive surface (e.g., 551 in FIG. 5B) has a primary axis (e.g., 552 in FIG. 5B) that corresponds to a primary axis (e.g., 553 in FIG. 5B) on the display (e.g., 550). In accordance with these embodiments, the device detects contacts (e.g., 560 and 562 in FIG. 5B) with the touch-sensitive surface 551 at locations that correspond to respective locations on the display (e.g., in FIG. 5B, contact 560 corresponds to 568 and contact 562 corresponds to 570). In this way, user inputs (e.g., contacts 560 and 562, and movements thereof) detected by the device on the touch-sensitive surface (e.g., 551 in FIG. 5B) are used by the device to manipulate the user interface on the display (e.g., 550 in FIG. 5B) of the multifunction device when the touch-sensitive surface is separate from the display. It should be understood that similar methods are, optionally, used for other user interfaces described herein.

[0191] Additionally, while the following examples are given primarily with reference to finger inputs (e.g., finger contacts, finger tap gestures, finger swipe gestures), it should be understood that, in some embodiments, one or more of the finger inputs are replaced with input from another input device (e.g., a mouse-based input or stylus input). For example, a swipe gesture is, optionally, replaced with a mouse click (e.g., instead of a contact) followed by movement of the cursor along the path of the swipe (e.g., instead of movement of the contact). As another example, a tap gesture is, optionally, replaced with a mouse click while the cursor is located over the location of the tap gesture (e.g., instead of detection of the contact followed by ceasing to detect the contact). Similarly, when multiple user inputs are simultaneously detected, it should be understood that multiple computer mice are, optionally, used simultaneously, or a mouse and finger contacts are, optionally, used simultaneously.

[0192] FIG. 6A illustrates exemplary personal electronic device 600. Device 600 includes body 602. In some embodiments, device 600 can include some or all of the features described with respect to devices 200 and 400 (e.g., FIGS. 2A-4B). In some embodiments, device 600 has touch-sensitive display screen 604, hereafter touch screen 604. Alternatively, or in addition to touch screen 604, device 600 has a display and a touch-sensitive surface. As with devices 200 and 400, in some embodiments, touch screen 604 (or the touch-sensitive surface) may have one or more intensity sensors for detecting intensity of contacts (e.g., touches) being applied. The one or more intensity sensors of touch screen 604 (or the touch-sensitive surface) can provide output data that represents the intensity of touches. The user interface of device 600 can respond to touches based on their

intensity, meaning that touches of different intensities can invoke different user interface operations on device **600**.

[0193] Techniques for detecting and processing touch intensity may be found, for example, in related applications: International Patent Application Serial No. PCT/US2013/040061, titled “Device, Method, and Graphical User Interface for Displaying User Interface Objects Corresponding to an Application,” filed May 8, 2013, and International Patent Application Serial No. PCT/US2013/069483, titled “Device, Method, and Graphical User Interface for Transitioning Between Touch Input to Display Output Relationships,” filed Nov. 11, 2013, each of which is hereby incorporated by reference in their entirety.

[0194] In some embodiments, device **600** has one or more input mechanisms **606** and **608**. Input mechanisms **606** and **608**, if included, can be physical. Examples of physical input mechanisms include push buttons and rotatable mechanisms. In some embodiments, device **600** has one or more attachment mechanisms. Such attachment mechanisms, if included, can permit attachment of device **600** with, for example, hats, eyewear, earrings, necklaces, shirts, jackets, bracelets, watch straps, chains, trousers, belts, shoes, purses, backpacks, and so forth. These attachment mechanisms may permit device **600** to be worn by a user.

[0195] FIG. 6B depicts exemplary personal electronic device **600**. In some embodiments, device **600** can include some or all of the components described with respect to FIGS. 2A, 2B, and 4. Device **600** has bus **612** that operatively couples I/O section **614** with one or more computer processors **616** and memory **618**. I/O section **614** can be connected to display **604**, which can have touch-sensitive component **622** and, optionally, touch-intensity sensitive component **624**. In addition, I/O section **614** can be connected with communication unit **630** for receiving application and operating system data, using Wi-Fi, Bluetooth, near field communication (NFC), cellular, and/or other wireless communication techniques. Device **600** can include input mechanisms **606** and/or **608**. Input mechanism **606** may be a rotatable input device or a depressible and rotatable input device, for example. Input mechanism **608** may be a button, in some examples.

[0196] Input mechanism **608** may be a microphone, in some examples. Personal electronic device **600** can include various sensors, such as GPS sensor **632**, accelerometer **634**, directional sensor **640** (e.g., compass), gyroscope **636**, motion sensor **638**, and/or a combination thereof, all of which can be operatively connected to I/O section **614**.

[0197] Memory **618** of personal electronic device **600** can be a non-transitory computer-readable storage medium, for storing computer-executable instructions, which, when executed by one or more computer processors **616**, for example, can cause the computer processors to perform the techniques described below, including processes **800-900** (FIGS. 8-9). The computer-executable instructions can also be stored and/or transported within any non-transitory computer-readable storage medium for use by or in connection with an instruction execution system, apparatus, or device, such as a computer-based system, processor-containing system, or other system that can fetch the instructions from the instruction execution system, apparatus, or device and execute the instructions. Personal electronic device **600** is not limited to the components and configuration of FIG. 6B, but can include other or additional components in multiple configurations.

[0198] As used here, the term “affordance” refers to a user-interactive graphical user interface object that may be displayed on the display screen of devices **200**, **400**, and/or **600** (FIGS. 2, 4, and 6). For example, an image (e.g., icon), a button, and text (e.g., link) may each constitute an affordance.

[0199] As used herein, the term “focus selector” refers to an input element that indicates a current part of a user interface with which a user is interacting. In some implementations that include a cursor or other location marker, the cursor acts as a “focus selector” so that when an input (e.g., a press input) is detected on a touch-sensitive surface (e.g., touchpad **455** in FIG. 4 or touch-sensitive surface **551** in FIG. 5B) while the cursor is over a particular user interface element (e.g., a button, window, slider or other user interface element), the particular user interface element is adjusted in accordance with the detected input. In some implementations that include a touch screen display (e.g., touch-sensitive display system **212** in FIG. 2A or touch screen **212** in FIG. 5A) that enables direct interaction with user interface elements on the touch screen display, a detected contact on the touch screen acts as a “focus selector” so that when an input (e.g., a press input by the contact) is detected on the touch screen display at a location of a particular user interface element (e.g., a button, window, slider, or other user interface element), the particular user interface element is adjusted in accordance with the detected input. In some implementations, focus is moved from one region of a user interface to another region of the user interface without corresponding movement of a cursor or movement of a contact on a touch screen display (e.g., by using a tab key or arrow keys to move focus from one button to another button); in these implementations, the focus selector moves in accordance with movement of focus between different regions of the user interface. Without regard to the specific form taken by the focus selector, the focus selector is generally the user interface element (or contact on a touch screen display) that is controlled by the user so as to communicate the user’s intended interaction with the user interface (e.g., by indicating, to the device, the element of the user interface with which the user is intending to interact). For example, the location of a focus selector (e.g., a cursor, a contact, or a selection box) over a respective button while a press input is detected on the touch-sensitive surface (e.g., a touchpad or touch screen) will indicate that the user is intending to activate the respective button (as opposed to other user interface elements shown on a display of the device).

[0200] As used in the specification and claims, the term “characteristic intensity” of a contact refers to a characteristic of the contact based on one or more intensities of the contact. In some embodiments, the characteristic intensity is based on multiple intensity samples. The characteristic intensity is, optionally, based on a predefined number of intensity samples, or a set of intensity samples collected during a predetermined time period (e.g., 0.05, 0.1, 0.2, 0.5, 1, 2, 5, 10 seconds) relative to a predefined event (e.g., after detecting the contact, prior to detecting liftoff of the contact, before or after detecting a start of movement of the contact, prior to detecting an end of the contact, before or after detecting an increase in intensity of the contact, and/or before or after detecting a decrease in intensity of the contact). A characteristic intensity of a contact is, optionally, based on one or more of: a maximum value of the intensities

of the contact, a mean value of the intensities of the contact, an average value of the intensities of the contact, a top 10 percentile value of the intensities of the contact, a value at the half maximum of the intensities of the contact, a value at the 90 percent maximum of the intensities of the contact, or the like. In some embodiments, the duration of the contact is used in determining the characteristic intensity (e.g., when the characteristic intensity is an average of the intensity of the contact over time). In some embodiments, the characteristic intensity is compared to a set of one or more intensity thresholds to determine whether an operation has been performed by a user. For example, the set of one or more intensity thresholds may include a first intensity threshold and a second intensity threshold. In this example, a contact with a characteristic intensity that does not exceed the first threshold results in a first operation, a contact with a characteristic intensity that exceeds the first intensity threshold and does not exceed the second intensity threshold results in a second operation, and a contact with a characteristic intensity that exceeds the second threshold results in a third operation. In some embodiments, a comparison between the characteristic intensity and one or more thresholds is used to determine whether or not to perform one or more operations (e.g., whether to perform a respective operation or forgo performing the respective operation) rather than being used to determine whether to perform a first operation or a second operation.

[0201] In some embodiments, a portion of a gesture is identified for purposes of determining a characteristic intensity. For example, a touch-sensitive surface may receive a continuous swipe contact transitioning from a start location and reaching an end location, at which point the intensity of the contact increases. In this example, the characteristic intensity of the contact at the end location may be based on only a portion of the continuous swipe contact, and not the entire swipe contact (e.g., only the portion of the swipe contact at the end location). In some embodiments, a smoothing algorithm may be applied to the intensities of the swipe contact prior to determining the characteristic intensity of the contact. For example, the smoothing algorithm optionally includes one or more of: an unweighted sliding-average smoothing algorithm, a triangular smoothing algorithm, a median filter smoothing algorithm, and/or an exponential smoothing algorithm. In some circumstances, these smoothing algorithms eliminate narrow spikes or dips in the intensities of the swipe contact for purposes of determining a characteristic intensity.

[0202] The intensity of a contact on the touch-sensitive surface may be characterized relative to one or more intensity thresholds, such as a contact-detection intensity threshold, a light press intensity threshold, a deep press intensity threshold, and/or one or more other intensity thresholds. In some embodiments, the light press intensity threshold corresponds to an intensity at which the device will perform operations typically associated with clicking a button of a physical mouse or a trackpad. In some embodiments, the deep press intensity threshold corresponds to an intensity at which the device will perform operations that are different from operations typically associated with clicking a button of a physical mouse or a trackpad. In some embodiments, when a contact is detected with a characteristic intensity below the light press intensity threshold (e.g., and above a nominal contact-detection intensity threshold below which the contact is no longer detected), the device will move a

focus selector in accordance with movement of the contact on the touch-sensitive surface without performing an operation associated with the light press intensity threshold or the deep press intensity threshold. Generally, unless otherwise stated, these intensity thresholds are consistent between different sets of user interface figures.

[0203] An increase of characteristic intensity of the contact from an intensity below the light press intensity threshold to an intensity between the light press intensity threshold and the deep press intensity threshold is sometimes referred to as a “light press” input. An increase of characteristic intensity of the contact from an intensity below the deep press intensity threshold to an intensity above the deep press intensity threshold is sometimes referred to as a “deep press” input. An increase of characteristic intensity of the contact from an intensity below the contact-detection intensity threshold to an intensity between the contact-detection intensity threshold and the light press intensity threshold is sometimes referred to as detecting the contact on the touch-surface. A decrease of characteristic intensity of the contact from an intensity above the contact-detection intensity threshold to an intensity below the contact-detection intensity threshold is sometimes referred to as detecting liftoff of the contact from the touch-surface. In some embodiments, the contact-detection intensity threshold is zero. In some embodiments, the contact-detection intensity threshold is greater than zero.

[0204] In some embodiments described herein, one or more operations are performed in response to detecting a gesture that includes a respective press input or in response to detecting the respective press input performed with a respective contact (or a plurality of contacts), where the respective press input is detected based at least in part on detecting an increase in intensity of the contact (or plurality of contacts) above a press-input intensity threshold. In some embodiments, the respective operation is performed in response to detecting the increase in intensity of the respective contact above the press-input intensity threshold (e.g., a “down stroke” of the respective press input). In some embodiments, the press input includes an increase in intensity of the respective contact above the press-input intensity threshold and a subsequent decrease in intensity of the contact below the press-input intensity threshold, and the respective operation is performed in response to detecting the subsequent decrease in intensity of the respective contact below the press-input threshold (e.g., an “up stroke” of the respective press input).

[0205] In some embodiments, the device employs intensity hysteresis to avoid accidental inputs sometimes termed “jitter,” where the device defines or selects a hysteresis intensity threshold with a predefined relationship to the press-input intensity threshold (e.g., the hysteresis intensity threshold is X intensity units lower than the press-input intensity threshold or the hysteresis intensity threshold is 75%, 90%, or some reasonable proportion of the press-input intensity threshold). Thus, in some embodiments, the press input includes an increase in intensity of the respective contact above the press-input intensity threshold and a subsequent decrease in intensity of the contact below the hysteresis intensity threshold that corresponds to the press-input intensity threshold, and the respective operation is performed in response to detecting the subsequent decrease in intensity of the respective contact below the hysteresis intensity threshold (e.g., an “up stroke” of the respective

press input). Similarly, in some embodiments, the press input is detected only when the device detects an increase in intensity of the contact from an intensity at or below the hysteresis intensity threshold to an intensity at or above the press-input intensity threshold and, optionally, a subsequent decrease in intensity of the contact to an intensity at or below the hysteresis intensity, and the respective operation is performed in response to detecting the press input (e.g., the increase in intensity of the contact or the decrease in intensity of the contact, depending on the circumstances). [0206] For ease of explanation, the descriptions of operations performed in response to a press input associated with a press-input intensity threshold or in response to a gesture including the press input are, optionally, triggered in response to detecting either: an increase in intensity of a contact above the press-input intensity threshold, an increase in intensity of a contact from an intensity below the hysteresis intensity threshold to an intensity above the press-input intensity threshold, a decrease in intensity of the contact below the press-input intensity threshold, and/or a decrease in intensity of the contact below the hysteresis intensity threshold corresponding to the press-input intensity threshold. Additionally, in examples where an operation is described as being performed in response to detecting a decrease in intensity of a contact below the press-input intensity threshold, the operation is, optionally, performed in response to detecting a decrease in intensity of the contact below a hysteresis intensity threshold corresponding to, and lower than, the press-input intensity threshold.

3. Digital Assistant System

[0207] FIG. 7A illustrates a block diagram of digital assistant system 700 in accordance with various examples. In some examples, digital assistant system 700 can be implemented on a standalone computer system. In some examples, digital assistant system 700 can be distributed across multiple computers. In some examples, some of the modules and functions of the digital assistant can be divided into a server portion and a client portion, where the client portion resides on one or more user devices (e.g., devices 104, 122, 200, 400, or 600) and communicates with the server portion (e.g., server system 108) through one or more networks, e.g., as shown in FIG. 1. In some examples, digital assistant system 700 can be an implementation of server system 108 (and/or DA server 106) shown in FIG. 1. It should be noted that digital assistant system 700 is only one example of a digital assistant system, and that digital assistant system 700 can have more or fewer components than shown, may combine two or more components, or may have a different configuration or arrangement of the components. The various components shown in FIG. 7A can be implemented in hardware, software instructions for execution by one or more processors, firmware, including one or more signal processing and/or application specific integrated circuits, or a combination thereof.

[0208] Digital assistant system 700 can include memory 702, one or more processors 704, input/output (I/O) interface 706, and network communications interface 708. These components can communicate with one another over one or more communication buses or signal lines 710.

[0209] In some examples, memory 702 can include a non-transitory computer-readable medium, such as high-speed random access memory and/or a non-volatile computer-readable storage medium (e.g., one or more magnetic

disk storage devices, flash memory devices, or other non-volatile solid-state memory devices).

[0210] In some examples, I/O interface 706 can couple input/output devices 716 of digital assistant system 700, such as displays, keyboards, touch screens, and microphones, to user interface module 722. I/O interface 706, in conjunction with user interface module 722, can receive user inputs (e.g., voice input, keyboard inputs, touch inputs, etc.) and processes them accordingly. In some examples, e.g., when the digital assistant is implemented on a standalone user device, digital assistant system 700 can include any of the components and I/O communication interfaces described with respect to devices 200, 400, or 600 in FIGS. 2A, 4, 6A-6B, respectively. In some examples, digital assistant system 700 can represent the server portion of a digital assistant implementation, and can interact with the user through a client-side portion residing on a user device (e.g., devices 104, 200, 400, or 600).

[0211] In some examples, the network communications interface 708 can include wired communication port(s) 712 and/or wireless transmission and reception circuitry 714. The wired communication port(s) can receive and send communication signals via one or more wired interfaces, e.g., Ethernet, Universal Serial Bus (USB), FIREWIRE, etc. The wireless circuitry 714 can receive and send RF signals and/or optical signals from/to communications networks and other communications devices. The wireless communications can use any of a plurality of communications standards, protocols, and technologies, such as GSM, EDGE, CDMA, TDMA, Bluetooth, Wi-Fi, VoIP, Wi-MAX, or any other suitable communication protocol. Network communications interface 708 can enable communication between digital assistant system 700 with networks, such as the Internet, an intranet, and/or a wireless network, such as a cellular telephone network, a wireless local area network (LAN), and/or a metropolitan area network (MAN), and other devices.

[0212] In some examples, memory 702, or the computer-readable storage media of memory 702, can store programs, modules, instructions, and data structures including all or a subset of: operating system 718, communications module 720, user interface module 722, one or more applications 724, and digital assistant module 726. In particular, memory 702, or the computer-readable storage media of memory 702, can store instructions for performing processes 800, 900, described below. One or more processors 704 can execute these programs, modules, and instructions, and reads/writes from/to the data structures.

[0213] Operating system 718 (e.g., Darwin, RTXC, LINUX, UNIX, iOS, OS X, WINDOWS, or an embedded operating system such as VxWorks) can include various software components and/or drivers for controlling and managing general system tasks (e.g., memory management, storage device control, power management, etc.) and facilitates communications between various hardware, firmware, and software components.

[0214] Communications module 720 can facilitate communications between digital assistant system 700 with other devices over network communications interface 708. For example, communications module 720 can communicate with RF circuitry 208 of electronic devices such as devices 200, 400, and 600 shown in FIG. 2A, 4, 6A-6B, respectively. Communications module 720 can also include various com-

ponents for handling data received by wireless circuitry 714 and/or wired communications port 712.

[0215] User interface module 722 can receive commands and/or inputs from a user via I/O interface 706 (e.g., from a keyboard, touch screen, pointing device, controller, and/or microphone), and generate user interface objects on a display. User interface module 722 can also prepare and deliver outputs (e.g., speech, sound, animation, text, icons, vibrations, haptic feedback, light, etc.) to the user via the I/O interface 706 (e.g., through displays, audio channels, speakers, touch-pads, etc.).

[0216] Applications 724 can include programs and/or modules that are configured to be executed by one or more processors 704. For example, if the digital assistant system is implemented on a standalone user device, applications 724 can include user applications, such as games, a calendar application, a navigation application, or an email application. If digital assistant system 700 is implemented on a server, applications 724 can include resource management applications, diagnostic applications, or scheduling applications, for example.

[0217] Memory 702 can also store digital assistant module 726 (or the server portion of a digital assistant). In some examples, digital assistant module 726 can include the following sub-modules, or a subset or superset thereof: input/output processing module 728, speech-to-text (STT) processing module 730, natural language processing module 732, dialogue flow processing module 734, task flow processing module 736, service processing module 738, and speech synthesis module 740. Each of these modules can have access to one or more of the following systems or data and models of the digital assistant module 726, or a subset or superset thereof: ontology 760, vocabulary index 744, user data 748, task flow models 754, service models 756, and ASR systems.

[0218] In some examples, using the processing modules, data, and models implemented in digital assistant module 726, the digital assistant can perform at least some of the following: converting speech input into text; identifying a user's intent expressed in a natural language input received from the user; actively eliciting and obtaining information needed to fully infer the user's intent (e.g., by disambiguating words, games, intentions, etc.); determining the task flow for fulfilling the inferred intent; and executing the task flow to fulfill the inferred intent.

[0219] In some examples, as shown in FIG. 7B, I/O processing module 728 can interact with the user through I/O devices 716 in FIG. 7A or with a user device (e.g., devices 104, 200, 400, or 600) through network communications interface 708 in FIG. 7A to obtain user input (e.g., a speech input) and to provide responses (e.g., as speech outputs) to the user input. I/O processing module 728 can optionally obtain contextual information associated with the user input from the user device, along with or shortly after the receipt of the user input. The contextual information can include user-specific data, vocabulary, and/or preferences relevant to the user input. In some examples, the contextual information also includes software and hardware states of the user device at the time the user request is received, and/or information related to the surrounding environment of the user at the time that the user request was received. In some examples, I/O processing module 728 can also send follow-up questions to, and receive answers from, the user regarding the user request. When a user request is received

by I/O processing module 728 and the user request can include speech input, I/O processing module 728 can forward the speech input to STT processing module 730 (or speech recognizer) for speech-to-text conversions.

[0220] STT processing module 730 can include one or more ASR systems. The one or more ASR systems can process the speech input that is received through I/O processing module 728 to produce a recognition result. Each ASR system can include a front-end speech pre-processor. The front-end speech pre-processor can extract representative features from the speech input. For example, the front-end speech pre-processor can perform a Fourier transform on the speech input to extract spectral features that characterize the speech input as a sequence of representative multi-dimensional vectors. Further, each ASR system can include one or more speech recognition models (e.g., acoustic models and/or language models) and can implement one or more speech recognition engines. Examples of speech recognition models can include Hidden Markov Models, Gaussian-Mixture Models, Deep Neural Network Models, n-gram language models, and other statistical models. Examples of speech recognition engines can include the dynamic time warping based engines and weighted finite-state transducers (WFST) based engines. The one or more speech recognition models and the one or more speech recognition engines can be used to process the extracted representative features of the front-end speech pre-processor to produce intermediate recognitions results (e.g., phonemes, phonemic strings, and sub-words), and ultimately, text recognition results (e.g., words, word strings, or sequence of tokens). In some examples, the speech input can be processed at least partially by a third-party service or on the user's device (e.g., device 104, 200, 400, or 600) to produce the recognition result. Once STT processing module 730 produces recognition results containing a text string (e.g., words, or sequence of words, or sequence of tokens), the recognition result can be passed to natural language processing module 732 for intent deduction.

[0221] More details on the speech-to-text processing are described in U.S. Utility application Ser. No. 13/236,942 for "Consolidating Speech Recognition Results," filed on Sep. 20, 2011, the entire disclosure of which is incorporated herein by reference.

[0222] In some examples, STT processing module 730 can include and/or access a vocabulary of recognizable words via phonetic alphabet conversion module 731. Each vocabulary word can be associated with one or more candidate pronunciations of the word represented in a speech recognition phonetic alphabet. In particular, the vocabulary of recognizable words can include a word that is associated with a plurality of candidate pronunciations. For example, the vocabulary may include the word "tomato" that is associated with the candidate pronunciations of /tə'meɪtoʊ/ and /tə'matoo/. Further, vocabulary words can be associated with custom candidate pronunciations that are based on previous speech inputs from the user. Such custom candidate pronunciations can be stored in STT processing module 730 and can be associated with a particular user via the user's profile on the device. In some examples, the candidate pronunciations for words can be determined based on the spelling of the word and one or more linguistic and/or phonetic rules. In some examples, the candidate pronunciations can be manually generated, e.g., based on known canonical pronunciations.

[0223] In some examples, the candidate pronunciations can be ranked based on the commonness of the candidate pronunciation. For example, the candidate pronunciation /tə'meɪrʊ/ can be ranked higher than /tə'mætʊ/, because the former is a more commonly used pronunciation (e.g., among all users, for users in a particular geographical region, or for any other appropriate subset of users). In some examples, candidate pronunciations can be ranked based on whether the candidate pronunciation is a custom candidate pronunciation associated with the user. For example, custom candidate pronunciations can be ranked higher than canonical candidate pronunciations. This can be useful for recognizing proper nouns having a unique pronunciation that deviates from canonical pronunciation. In some examples, candidate pronunciations can be associated with one or more speech characteristics, such as geographic origin, nationality, or ethnicity. For example, the candidate pronunciation /tə'meɪrʊ/ can be associated with the United States, whereas the candidate pronunciation /tə'mætʊ/ can be associated with Great Britain. Further, the rank of the candidate pronunciation can be based on one or more characteristics (e.g., geographic origin, nationality, ethnicity, etc.) of the user stored in the user's profile on the device. For example, it can be determined from the user's profile that the user is associated with the United States. Based on the user being associated with the United States, the candidate pronunciation /tə'meɪrʊ/ (associated with the United States) can be ranked higher than the candidate pronunciation /tə'mætʊ/ (associated with Great Britain). In some examples, one of the ranked candidate pronunciations can be selected as a predicted pronunciation (e.g., the most likely pronunciation).

[0224] When a speech input is received, STT processing module 730 can be used to determine the phonemes corresponding to the speech input (e.g., using an acoustic model), and then attempt to determine words that match the phonemes (e.g., using a language model). For example, if STT processing module 730 can first identify the sequence of phonemes /tə'meɪrʊ/ corresponding to a portion of the speech input, it can then determine, based on vocabulary index 744, that this sequence corresponds to the word "tomato."

[0225] In some examples, STT processing module 730 can use approximate matching techniques to determine words in an utterance. Thus, for example, the STT processing module 730 can determine that the sequence of phonemes /tə'meɪrʊ/ corresponds to the word "tomato," even if that particular sequence of phonemes is not one of the candidate sequence of phonemes for that word.

[0226] Natural language processing module 732 ("natural language processor") of the digital assistant can take the sequence of words or tokens ("token sequence") generated by STT processing module 730, and attempt to associate the token sequence with one or more "actionable intents" recognized by the digital assistant. An "actionable intent" can represent a task that can be performed by the digital assistant, and can have an associated task flow implemented in task flow models 754. The associated task flow can be a series of programmed actions and steps that the digital assistant takes in order to perform the task. The scope of a digital assistant's capabilities can be dependent on the number and variety of task flows that have been implemented and stored in task flow models 754, or in other words, on the number and variety of "actionable intents" that

the digital assistant recognizes. The effectiveness of the digital assistant, however, can also be dependent on the assistant's ability to infer the correct "actionable intent(s)" from the user request expressed in natural language.

[0227] In some examples, in addition to the sequence of words or tokens obtained from STT processing module 730, natural language processing module 732 can also receive contextual information associated with the user request, e.g., from I/O processing module 728. The natural language processing module 732 can optionally use the contextual information to clarify, supplement, and/or further define the information contained in the token sequence received from STT processing module 730. The contextual information can include, for example, user preferences, hardware, and/or software states of the user device, sensor information collected before, during, or shortly after the user request, prior interactions (e.g., dialogue) between the digital assistant and the user, and the like. As described herein, contextual information can be dynamic, and can change with time, location, content of the dialogue, and other factors.

[0228] In some examples, the natural language processing can be based on, e.g., ontology 760. Ontology 760 can be a hierarchical structure containing many nodes, each node representing either an "actionable intent" or a "property" relevant to one or more of the "actionable intents" or other "properties." As noted above, an "actionable intent" can represent a task that the digital assistant is capable of performing, i.e., it is "actionable" or can be acted on. A "property" can represent a parameter associated with an actionable intent or a sub-aspect of another property. A linkage between an actionable intent node and a property node in ontology 760 can define how a parameter represented by the property node pertains to the task represented by the actionable intent node.

[0229] In some examples, ontology 760 can be made up of actionable intent nodes and property nodes. Within ontology 760, each actionable intent node can be linked to one or more property nodes either directly or through one or more intermediate property nodes. Similarly, each property node can be linked to one or more actionable intent nodes either directly or through one or more intermediate property nodes. For example, as shown in FIG. 7C, ontology 760 can include a "restaurant reservation" node (i.e., an actionable intent node). Property nodes "restaurant," "date/time" (for the reservation), and "party size" can each be directly linked to the actionable intent node (i.e., the "restaurant reservation" node).

[0230] In addition, property nodes "cuisine," "price range," "phone number," and "location" can be sub-nodes of the property node "restaurant," and can each be linked to the "restaurant reservation" node (i.e., the actionable intent node) through the intermediate property node "restaurant." For another example, as shown in FIG. 7C, ontology 760 can also include a "set reminder" node (i.e., another actionable intent node). Property nodes "date/time" (for setting the reminder) and "subject" (for the reminder) can each be linked to the "set reminder" node. Since the property "date/time" can be relevant to both the task of making a restaurant reservation and the task of setting a reminder, the property node "date/time" can be linked to both the "restaurant reservation" node and the "set reminder" node in ontology 760.

[0231] An actionable intent node, along with its linked concept nodes, can be described as a "domain." In the

present discussion, each domain can be associated with a respective actionable intent, and refers to the group of nodes (and the relationships there between) associated with the particular actionable intent. For example, ontology 760 shown in FIG. 7C can include an example of restaurant reservation domain 762 and an example of reminder domain 764 within ontology 760. The restaurant reservation domain includes the actionable intent node “restaurant reservation,” property nodes “restaurant,” “date/time,” and “party size,” and sub-property nodes “cuisine,” “price range,” “phone number,” and “location.” Reminder domain 764 can include the actionable intent node “set reminder,” and property nodes “subject” and “date/time.” In some examples, ontology 760 can be made up of many domains. Each domain can share one or more property nodes with one or more other domains. For example, the “date/time” property node can be associated with many different domains (e.g., a scheduling domain, a travel reservation domain, a movie ticket domain, etc.), in addition to restaurant reservation domain 762 and reminder domain 764.

[0232] While FIG. 7C illustrates two example domains within ontology 760, other domains can include, for example, “find a movie,” “initiate a phone call,” “find directions,” “schedule a meeting,” “send a message,” and “provide an answer to a question,” “read a list,” “providing navigation instructions,” “provide instructions for a task” and so on. A “send a message” domain can be associated with a “send a message” actionable intent node, and may further include property nodes such as “recipient(s),” “message type,” and “message body.” The property node “recipient” can be further defined, for example, by the sub-property nodes such as “recipient name” and “message address.”

[0233] In some examples, ontology 760 can include all the domains (and hence actionable intents) that the digital assistant is capable of understanding and acting upon. In some examples, ontology 760 can be modified, such as by adding or removing entire domains or nodes, or by modifying relationships between the nodes within the ontology 760.

[0234] In some examples, nodes associated with multiple related actionable intents can be clustered under a “super domain” in ontology 760. For example, a “travel” super-domain can include a cluster of property nodes and actionable intent nodes related to travel. The actionable intent nodes related to travel can include “airline reservation,” “hotel reservation,” “car rental,” “get directions,” “find points of interest,” and so on. The actionable intent nodes under the same super domain (e.g., the “travel” super domain) can have many property nodes in common. For example, the actionable intent nodes for “airline reservation,” “hotel reservation,” “car rental,” “get directions,” and “find points of interest” can share one or more of the property nodes “start location,” “destination,” “departure date/time,” “arrival date/time,” and “party size.”

[0235] In some examples, each node in ontology 760 can be associated with a set of words and/or phrases that are relevant to the property or actionable intent represented by the node. The respective set of words and/or phrases associated with each node can be the so-called “vocabulary” associated with the node. The respective set of words and/or phrases associated with each node can be stored in vocabulary index 744 in association with the property or actionable intent represented by the node. For example, returning to FIG. 7B, the vocabulary associated with the node for the

property of “restaurant” can include words such as “food,” “drinks,” “cuisine,” “hungry,” “eat,” “pizza,” “fast food,” “meal,” and so on. For another example, the vocabulary associated with the node for the actionable intent of “initiate a phone call” can include words and phrases such as “call,” “phone,” “dial,” “ring,” “call this number,” “make a call to,” and so on. The vocabulary index 744 can optionally include words and phrases in different languages.

[0236] Natural language processing module 732 can receive the token sequence (e.g., a text string) from STT processing module 730, and determine what nodes are implicated by the words in the token sequence. In some examples, if a word or phrase in the token sequence is found to be associated with one or more nodes in ontology 760 (via vocabulary index 744), the word or phrase can “trigger” or “activate” those nodes. Based on the quantity and/or relative importance of the activated nodes, natural language processing module 732 can select one of the actionable intents as the task that the user intended the digital assistant to perform. In some examples, the domain that has the most “triggered” nodes can be selected. In some examples, the domain having the highest confidence value (e.g., based on the relative importance of its various triggered nodes) can be selected. In some examples, the domain can be selected based on a combination of the number and the importance of the triggered nodes. In some examples, additional factors are considered in selecting the node as well, such as whether the digital assistant has previously correctly interpreted a similar request from a user.

[0237] User data 748 can include user-specific information, such as user-specific vocabulary, user preferences, user address, user’s default and secondary languages, user’s contact list, and other short-term or long-term information for each user. In some examples, natural language processing module 732 can use the user-specific information to supplement the information contained in the user input to further define the user intent. For example, for a user request “invite my friends to my birthday party,” natural language processing module 732 can be able to access user data 748 to determine who the “friends” are and when and where the “birthday party” would be held, rather than requiring the user to provide such information explicitly in his/her request.

[0238] Other details of searching an ontology based on a token string is described in U.S. Utility application Ser. No. 12/341,743 for “Method and Apparatus for Searching Using An Active Ontology,” filed Dec. 22, 2008, the entire disclosure of which is incorporated herein by reference.

[0239] In some examples, once natural language processing module 732 identifies an actionable intent (or domain) based on the user request, natural language processing module 732 can generate a structured query to represent the identified actionable intent. In some examples, the structured query can include parameters for one or more nodes within the domain for the actionable intent, and at least some of the parameters are populated with the specific information and requirements specified in the user request. For example, the user may say “Make me a dinner reservation at a sushi place at 7.” In this case, natural language processing module 732 can be able to correctly identify the actionable intent to be “restaurant reservation” based on the user input. According to the ontology, a structured query for a “restaurant reservation” domain may include parameters such as {Cuisine}, {Time}, {Date}, {Party Size}, and the like. In some

examples, based on the speech input and the text derived from the speech input using STT processing module 730, natural language processing module 732 can generate a partial structured query for the restaurant reservation domain, where the partial structured query includes the parameters {Cuisine="Sushi"} and {Time="7 pm"}. However, in this example, the user's utterance contains insufficient information to complete the structured query associated with the domain. Therefore, other necessary parameters such as {Party Size} and {Date} may not be specified in the structured query based on the information currently available. In some examples, natural language processing module 732 can populate some parameters of the structured query with received contextual information. For example, in some examples, if the user requested a sushi restaurant "near me," natural language processing module 732 can populate a {location} parameter in the structured query with GPS coordinates from the user device.

[0240] In some examples, natural language processing module 732 can pass the generated structured query (including any completed parameters) to task flow processing module 736 ("task flow processor"). Task flow processing module 736 can be configured to receive the structured query from natural language processing module 732, complete the structured query, if necessary, and perform the actions required to "complete" the user's ultimate request. In some examples, the various procedures necessary to complete these tasks can be provided in task flow models 754. In some examples, task flow models 754 can include procedures for obtaining additional information from the user and task flows for performing actions associated with the actionable intent.

[0241] As described above, in order to complete a structured query, task flow processing module 736 may need to initiate additional dialogue with the user in order to obtain additional information, and/or disambiguate potentially ambiguous utterances. When such interactions are necessary, task flow processing module 736 can invoke dialogue flow processing module 734 to engage in a dialogue with the user. In some examples, dialogue flow processing module 734 can determine how (and/or when) to ask the user for the additional information and receives and processes the user responses. The questions can be provided to and answers can be received from the users through I/O processing module 728. In some examples, dialogue flow processing module 734 can present dialogue output to the user via audio and/or visual output, and receives input from the user via spoken or physical (e.g., clicking) responses. Continuing with the example above, when task flow processing module 736 invokes dialogue flow processing module 734 to determine the "party size" and "date" information for the structured query associated with the domain "restaurant reservation," dialogue flow processing module 734 can generate questions such as "For how many people?" and "On which day?" to pass to the user. Once answers are received from the user, dialogue flow processing module 734 can then populate the structured query with the missing information, or pass the information to task flow processing module 736 to complete the missing information from the structured query.

[0242] Once task flow processing module 736 has completed the structured query for an actionable intent, task flow processing module 736 can proceed to perform the ultimate task associated with the actionable intent. Accordingly, task flow processing module 736 can execute the steps and

instructions in the task flow model according to the specific parameters contained in the structured query. For example, the task flow model for the actionable intent of "restaurant reservation" can include steps and instructions for contacting a restaurant and actually requesting a reservation for a particular party size at a particular time. For example, using a structured query such as: {restaurant reservation, restaurant=ABC Café, date=3/12/2012, time=7 pm, party size=5}, task flow processing module 736 can perform the steps of: (1) logging onto a server of the ABC Café or a restaurant reservation system such as OPENTABLE®, (2) entering the date, time, and party size information in a form on the website, (3) submitting the form, and (4) making a calendar entry for the reservation in the user's calendar.

[0243] In some examples, task flow processing module 736 can employ the assistance of service processing module 738 ("service processing module") to complete a task requested in the user input or to provide an informational answer requested in the user input. For example, service processing module 738 can act on behalf of task flow processing module 736 to make a phone call, set a calendar entry, invoke a map search, invoke or interact with other user applications installed on the user device, and invoke or interact with third-party services (e.g., a restaurant reservation portal, a social networking website, a banking portal, etc.). In some examples, the protocols and application programming interfaces (API) required by each service can be specified by a respective service model among service models 756. Service processing module 738 can access the appropriate service model for a service and generate requests for the service in accordance with the protocols and APIs required by the service according to the service model.

[0244] For example, if a restaurant has enabled an online reservation service, the restaurant can submit a service model specifying the necessary parameters for making a reservation and the APIs for communicating the values of the necessary parameter to the online reservation service. When requested by task flow processing module 736, service processing module 738 can establish a network connection with the online reservation service using the web address stored in the service model, and send the necessary parameters of the reservation (e.g., time, date, party size) to the online reservation interface in a format according to the API of the online reservation service.

[0245] In some examples, natural language processing module 732, dialogue flow processing module 734, and task flow processing module 736 can be used collectively and iteratively to infer and define the user's intent, obtain information to further clarify and refine the user intent, and finally generate a response (i.e., an output to the user, or the completion of a task) to fulfill the user's intent. The generated response can be a dialogue response to the speech input that at least partially fulfills the user's intent. Further, in some examples, the generated response can be output as a speech output. In these examples, the generated response can be sent to speech synthesis module 740 (e.g., speech synthesizer) where it can be processed to synthesize the dialogue response in speech form. In yet other examples, the generated response can be data content relevant to satisfying a user request in the speech input.

[0246] Speech synthesis module 740 can be configured to synthesize speech outputs for presentation to the user. Speech synthesis module 740 synthesizes speech outputs based on text provided by the digital assistant. For example,

the generated dialogue response can be in the form of a text string. Speech synthesis module **740** can convert the text string to an audible speech output. Speech synthesis module **740** can use any appropriate speech synthesis technique in order to generate speech outputs from text, including, but not limited, to concatenative synthesis, unit selection synthesis, diphone synthesis, domain-specific synthesis, formant synthesis, articulatory synthesis, hidden Markov model (HMM) based synthesis, and sinewave synthesis. In some examples, speech synthesis module **740** can be configured to synthesize individual words based on phonemic strings corresponding to the words. For example, a phonemic string can be associated with a word in the generated dialogue response. The phonemic string can be stored in metadata associated with the word. Speech synthesis model **740** can be configured to directly process the phonemic string in the metadata to synthesize the word in speech form.

[0247] In some examples, instead of (or in addition to) using speech synthesis module **740**, speech synthesis can be performed on a remote device (e.g., the server system **108**), and the synthesized speech can be sent to the user device for output to the user. For example, this can occur in some implementations where outputs for a digital assistant are generated at a server system. And because server systems generally have more processing power or resources than a user device, it can be possible to obtain higher quality speech outputs than would be practical with client-side synthesis.

[0248] Additional details on digital assistants can be found in the U.S. Utility application Ser. No. 12/987,982, entitled “Intelligent Automated Assistant,” filed Jan. 10, 2011, and U.S. Utility application Ser. No. 13/251,088, entitled “Generating and Processing Task Items That Represent Tasks to Perform,” filed Sep. 30, 2011, the entire disclosures of which are incorporated herein by reference.

4. Processes for Operating a Digital Assistant

[0249] FIG. 8 illustrates a flow diagram of a process **800** for operating a digital assistant in accordance with some embodiments. Process **800** is performed, for example, using one or more electronic devices (e.g., devices **104**, **108**, **200**, **400**, or **600**) implementing a digital assistant. In some examples, the process **800** is performed using a client-server system (e.g., system **100**), and the blocks of the process **800** can be divided up in any manner between the server (e.g., DA server **106**) and a client device. Thus, while portions of the process **800** are described herein as being performed by particular devices of a client-server system, it will be appreciated that the process **800** is not so limited. In other examples, the process **800** is performed using only a client device (e.g., user device **104**). In process **800**, some blocks are, optionally, combined, the order of some blocks is, optionally, changed, and some blocks are, optionally, omitted. In some examples, additional steps may be performed in combination with the process **800**.

[0250] At block **805**, a natural-language user input is received by a user device, such as the user device **104** of FIG. 1. The natural-language input is, for example, a speech input or a text input. In some examples, the natural-language input may include a request for the user device, and/or another device, to perform a task. For instance, in the example “Send a car to 1200 Main Street,” the natural-language input may include a request for a user device to reserve a car using a ride-booking service. In some

examples, the natural-language input may further specify one or more parameters for the requested task. “1200 Main Street”, for instance, specifies a pick-up location for the car reservation. In the example, “Order my usual from Domino’s,” the natural-language input may include a request for a user device to order food from the pizza chain Domino’s. “my usual” may further specify, contextually, what food is to be ordered.

[0251] At block **810**, an intent, and optionally, one or more parameters associated with the intent, are identified. The intent and parameters may be derived from the natural-language user input, for example. As described, intents may correspond to tasks requested by a user. Accordingly, identifying (e.g., determining) an intent may include identifying a task specified in the natural-language user input and/or inferring an intent corresponding to a requested task based on language and/or context of the natural-language user input. Intents may correspond to any type of tasks performed by a user device and, in particular, may correspond to tasks performed by one or more applications of a user device, as described in further detail below.

[0252] In some examples, intents are associated with (e.g., included in) one or more domains (e.g., intent categories, set of intents). Each domain may include a particular class of intents, allowing for intents to be intuitively grouped. For example, intents to reserve a car, cancel a car reservation, and/or any other intents directed to tasks commonly associated with ride-booking may be included in a ride-booking domain. In another example, intents to check-in to a flight, cancel a flight, reschedule a flight, retrieve flight information, and/or any other intents directed to tasks commonly associated with air travel may be included in an air travel domain. In another example, intents to provide directions, retrieve traffic information, and/or any other intents directed to tasks commonly associated with navigation may be included in a navigation domain. In yet another example, intents to send payments, receive payments, and/or any other intents directed to tasks commonly associated with financial transactions may be included in a financial transactions domain.

[0253] Identifying parameters may include identifying portions of the natural-language input that specify a manner in which a task corresponding to the intent is to be performed. Parameters may, for instance, specify locations (e.g., addresses or places of interest), times, dates, contacts, types, text (e.g., to be inserted into an email or message), quantities (e.g., distance, money) and, in some instances, names of software applications to perform the task. Parameters may further specify other conditions for tasks, examples of which are described herein.

[0254] Parameters may be identified, for instance, using one or more detectors. Each of the detectors may be configured to parse the natural-language user input (e.g., a textual representation of the natural-language user input) and identify one or more respective data types. By way of example, a first detector may be configured to identify user contacts and a second detector may be configured to identify addresses. Other detectors may identify data types including, but not limited to, phone numbers, names, persons of interest, places of interest, URLs, times, flight numbers, package tracking numbers, and dates.

[0255] In some examples, words of a custom vocabulary may be identified as parameters. For instance, one or more detectors may be configured to identify custom vocabulary

of one or more applications, respectively. Custom vocabulary of an application may include a name of the application (e.g., Uber, Lyft, Instagram, Flickr, WeChat, WhatsApp, LINE, Viber) and/or may include other terms uniquely associated with the application (e.g., UberX, DM, Lyftline, ZipCar).

[0256] In some examples, one or more applications may register with an application registration service. The service may be hosted by or otherwise accessible to the server 108 and/or the user device 104. Registering in this manner may include specifying one or more custom vocabulary terms associated with the application and, optionally, one or more language models for the custom vocabulary terms. The language models may, for instance, provide one or more pronunciations for each of the custom vocabulary terms. Language models provided in this manner may thereafter be used to help identify use of such custom terms during analysis of natural-language user inputs. In some examples, custom vocabulary may be included in the vocabulary index 744 (FIG. 7B).

[0257] In some examples, one or more parameters may be inferred from the natural-language user input. In the example “Get me a ride to the stadium,” for instance, it may be inferred that one parameter associated with the intent is a location of the user. In another example, “Pay John back for the meal,” it may be inferred that one parameter associated with the intent is an amount of money.

[0258] In some examples, an intent is identified based on the natural-language user input, and parameters associated with the intent are thereafter identified. Further, in some examples, parameters not associated with the intent are not identified. By way of example, an intent corresponding to directions (e.g., driving directions) may be associated with parameters specifying one or more locations (e.g., an origin and/or a destination), and/or a mode of transit. Consider, for instance, the example “Get me real-time driving directions to 1200 Main Street.” In this instance, the identified intent corresponds to a task for providing directions, “driving” is a parameter specifying a mode of transit, and “1200 Main Street” is a parameter specifying a location. The portion of the user input “real-time” is not a parameter associated with the identified intent. Accordingly, while “real-time” may be a valid parameter for some intents, “real-time” will not be identified as a parameter during operation.

[0259] In other examples, one or more parameters first may be identified and the intent may be identified based on the identified one or more parameters. In yet other examples, the intent and parameters associated with the intent may be identified concurrently.

[0260] In some examples, intents and parameters of a natural-language user input are identified by the user device, such as the user device 104 of FIG. 1. In other examples, the user device provides the natural-language user input (or a representation thereof) to a server, such as the server 108 of FIG. 1, and the server identifies (e.g., determines) intents and parameters of the natural-language user input, as described. Thereafter the server provides (e.g., transmits) the identified intents and parameters to the user device.

[0261] Optionally, once an intent and any parameters have been identified, the user device confirms identified intents and/or parameters with a user of the user device, and in some instances, confirms inferred parameters with the user of the user device. Confirming in this manner may include prompting a user to confirm the identified intent and all identified parameters associated with the intent in response to a natural-language query. By way of example, in response to the user input “Get me a ride to the airport,” the user device may provide a natural-language query “Do you want a ride to the airport from your current location?” The natural-language query provided by the user device may be provided to a user as text, using a touch-sensitive display of the user device, and/or may be provided to a user as speech using an audio output component of the user device (e.g., speaker 211 of FIG. 2). The user may respond to the natural-language query, for instance, by providing a natural-language user input to the user device.

[0262] Optionally, the user device confirms individual parameters. In some examples, this may include prompting a user to confirm one or more parameters. By way of example, in response to the user input “Get me a ride to the station,” the user device may provide a natural language query “Do you mean Penn Station?” As another example, in response to the user input “Pay John \$5.00,” the user device may provide a natural language query “Do you mean John Smith?” Natural-language queries provided by the user device may be provided to a user as text, using a touch-sensitive display of the user device and/or may be provided to a user as speech using an audio output component of the user device (e.g., speaker 211 of FIG. 2). The user may respond to the natural-language query, for instance, by providing a natural-language user input to the user device.

[0263] In some examples, one or more parameters are contextual. Accordingly, the user device may determine (e.g., resolve) one or more parameters based on context information. Context information may be context information of the user device (or any data stored therein) and/or context information of a user of the user device. By way of example, a natural-language user input may recite “Get me a car at my house.” Because “my house” is a contextual parameter and does not specify an actual location, the user device may determine a location of the user device and identify the determined location as a parameter (i.e., in lieu of “my house”). As another example, a natural-language user input may recite “Call him back.” Because “him” is a contextual parameter and does not specify a specific contact, the user device may determine a contact intended by “him” and pass the contact as a parameter (i.e., in lieu of “him”).

[0264] In some examples, identified intents and parameters are implemented as an intent object. When implemented as such, each intent object is an object (e.g., data structure, programming object) and corresponds to a respective intent. Each intent object may include one or more fields (e.g., instance variables) corresponding to one or more parameters, respectively. For example, an intent object corresponding to a ride-booking intent may be generated (e.g., instantiated) like the following pseudocode:

```
public final class RideBookingIntent {
    public final static String INTENT_TYPE = "RideBookingIntent";
    public final parameter<Location> pickupLocation;
```

-continued

```
public final parameter<Location> dropOffLocation;
public final parameter<dataString> vehicleType;
public final parameter<Integer> partySize;
public final parameter<dateTime> pickupTime;
private RideBookingIntent(Location p, Location d, dataString v, Integer p, dateTime pt) {
    pickupLocation = p;
    dropOffLocation = d;
    vehicleType = v;
    partySize = p;
    pickupTime = pt; } }
```

As will be appreciated by one of ordinary skill, the pseudo-code above is exemplary, and an intent object may be implemented in other manners.

[0265] By implementing an intent as an intent object, intents may be language agnostic. As described, intents may be derived from natural-language user inputs. A same intent, therefore, may be derived from natural-language inputs provided in any number of spoken languages. By way of example, the English natural-language user input “Send an Uber to 1200 Park Avenue” and the German natural-language user input “Senden Sie eine Uber auf 1200 Park Avenue” would each result in a same intent being identified (and result in a same intent object).

[0266] At block 815, a software application associated with the intent may be identified (e.g., selected). Generally, this may include identifying one or more software applications configured to perform a task corresponding to the intent.

[0267] In some examples, identifying software applications may include determining one or more domains corresponding to an intent and identifying applications corresponding to the domain(s). As described, one or more software applications may be registered with an application registration service. Registering in this manner may include specifying which domains (e.g., ride-booking domain, air travel domain, navigation domain) correspond to the software application. An application corresponding to a domain may support each of the intents specified by the domain or may support only some intents specified by the domain. In some examples, applications may be registered with individual intents, and identifying applications may include identifying applications corresponding to the identified intent.

[0268] In some examples, applications are identified according to the identified parameters. By way of example, based on a user input “Get me a black car to the airport,” an intent to reserve a car using a ride-booking service and a parameter specifying a type of car (i.e. “black car”) may be identified. Although several available applications may be configured to reserve a car generally, only those applications configured to reserve a “black car” may be identified. As another example, based on a user input “Send a message to Sam saying ‘Hello!’” Although several applications may be configured to send a message, only those applications having contact information for a contact Sam may be identified.

[0269] In some examples, only applications installed on and/or accessible to the user device are identified. For instance, although several available applications may be configured to perform a task, only those accessible to the user device are identified. Accessible applications include applications residing and/or installed on the user device and

further include applications that may be remotely accessed by the user device, for instance on one or more other devices.

[0270] Thus, in at least some embodiments, identified applications are applications configured to perform a task according to identified parameters and accessible to the user device. In some examples, multiple applications may satisfy these criteria, yet the user device may wish to identify fewer applications, or a single application. Accordingly, applications may further be identified based on previous use of the applications by the user device. In some examples, for a given intent, an application most recently used to perform a task corresponding to the intent is identified. Consider the user input “Place a call to Rob” in which an intent (i.e., place a call) and a parameter (i.e. “Rob”) may be identified. In this example, the application last used to place a call is identified. In other examples, an application most commonly used to perform tasks corresponding to intent is identified. For the same example, an application most commonly used to place calls is identified. In some examples, applications are further selected in accordance with one or more parameters. For instance, an application recently used to place a call to the contact Rob is identified, or an application most commonly used to place a call to the contact Rob is identified. In some examples, a default application may be specified for one or more particular tasks and/or parameters, for instance, by a user or a digital assistant. A user may specify, for instance, that a first application is to be used when calling a first contact and a second application is to be used when calling a second contact.

[0271] As described, in some instances, a natural-language user input may include custom vocabulary that may be identified as one or more parameters. In some examples, such custom vocabulary includes application names, and accordingly an application may be identified based on presence of custom vocabulary in the input. A natural-language user input may, for instance, recite “Call Rob using Skype.” In response, the software application Skype may be the identified software application. In another example, a natural-language input may recite “Play The Beatles on Spotify,” and in response, the software application Spotify may be the identified software application.

[0272] Custom vocabulary may further include terms uniquely associated with an application. Accordingly, such terms may be identified as parameters and used to identify an application. In the example “Get me an UberX,” “UberX” is a term of a custom vocabulary for the software application Uber and as a result Uber is identified as the software application. In the example “Tweet I hope the Sharks win the cup,” “Tweet” is a term of a custom vocabulary for the software application Twitter, and as a result Twitter is identified as the software application.

[0273] In some examples, no application configured to perform the task according to the identified parameters may be accessible to the user device. As a result, the user device may access (e.g., download and/or install) an application configured to perform the task according to the identified parameters. In some examples, the user device may identify a plurality of software applications and provide the user with a list of software applications. The user may select one or more of the applications, and the user device may access the one or more selected applications.

[0274] At block 820, the intent and parameters are provided to the identified software application. In some examples, the intent and parameters are provided to the software application as an intent object.

[0275] In some examples, the intent and parameters may be selectively provided to the software application based on a state of the user device. As an example, intents and parameters may be selectively provided to the user device based on whether the user device is in a lock state. In some examples, an application may be allowed to receive particular intents and parameters when the device is in a lock state. In other examples, an application may be allowed to receive particular intents and parameters when the user device is not in a lock state. Whether an application can receive a particular intent given a particular state of the user device may be specified by the software application, for instance during a registration process with an application registration service.

[0276] At block 825, the user device may receive one or more responses from the software application. In some examples, the user device receives a response for each parameter provided to the software application. Each of the responses may indicate, for instance, whether a parameter is valid or whether additional user input is required. If a response indicates that a parameter is valid, no further action is taken with respect to the parameter.

[0277] If a response provided by the software application does not indicate that a parameter is valid, the response may indicate that clarification of the parameter is required. For instance, a parameter may be improper (i.e., invalid) and the software application may request additional input from the user. As an example, consider the user input "Send a blue car to 1200 Main Street." Although the user has requested a blue car, the ride-booking application (e.g., Uber, Lyft) may not allow for the selection of a blue car (e.g., blue car may not be a supported parameter or the application may determine that no blue car is currently available). Accordingly, in the event the parameter is improper (e.g., the user has specified an invalid type of car), the application may request that a proper (e.g., valid) value for the parameter (e.g., type of car) be provided. With reference to FIG. 10A, for instance, based on the response from the software application indicating the parameter is improper, the user device may provide a natural-language query 1002 prompting the user to select a valid parameter. In the instant example, the user device provides a natural-language query asking the user to select a valid type of car. The natural-language query 1002 may be provided to a user as text using a touch-sensitive display of the user device and/or may be provided to a user as speech using an audio output component of the user device. As illustrated, in some examples, the user device may provide (e.g., display) one or more candidate parameters 1004 to the user for selection. In the instant example, candidate parameters 1004 include "Budget," "Black Car," "SUV," and

"Shared." Candidate parameters 1004 provided in this manner may be provided by the software application in some examples. The user device may select one of the candidate parameters by providing a touch-input and/or by providing a natural-language user input to the user device, and in response, the user device may provide the selected candidate parameter to the software application.

[0278] If a response provided by the software application does not indicate that a parameter is valid, the response may indicate that disambiguation of the parameter is required. As a result the user device may request input from the user to disambiguate the parameter. In the example user input "Call Tom," a parameter for a contact Tom may be provided to a software application configured to place calls. If, in determining contact information (e.g., a phone number) for the contact Tom, the software application determines that multiple contacts with the name "Tom" exist, the software application may request that the user specify which "Tom" was intended. As part of the request, the software application includes in the response a disambiguation list including a plurality of candidate parameters. The user device may provide a natural-language query asking the user to select a candidate parameter. Further, candidate parameters of the disambiguation list may be displayed to allow for user selection. The user may select one of the candidate parameters, for instance, by providing a touch-input and/or by providing a natural-language user input, and the user device may provide the selected candidate parameter to the software application.

[0279] Performing a task may require that one or more specific types of parameters are specified by a user. Yet in some examples, a natural-language user input may omit one or more required parameters. Accordingly, the software application may optionally provide one or more responses indicating which, if any, required parameters were not specified. Consider the user input "Send a car to 1200 Main Street." Although a car is generally requested by way of the user input, the ride-booking application identified based on the user input may require a selection of a particular type of car. Accordingly, in the event the parameter is missing (e.g., the user has not specified a type of car), the application may request that a recognized value for the parameter (e.g., type of car) be provided. With reference once again to FIG. 10A, the user device may provide a natural-language query 1002 prompting the user to specify a valid parameter. Thereafter the user may select a candidate parameter, for instance from a list of candidate parameters, and the selected parameter may be provided to the software application as described.

[0280] Once the software application has indicated that each parameter is valid and no additional information is required, at block 830, the user device may confirm intent with the software application. In particular, the user device may request a notification that, given the intent and the parameters associated with the intent, the software application can successfully perform a task corresponding to the intent.

[0281] Once the software application provides a notification indicating that the software application can perform the task, optionally, the user device confirms intent with the user. For example, the user device may provide the natural-language query "I can get an uberX at your location. Shall I request it?" The user may confirm or reject the intent by way of either a touch input or a natural-language input. In some examples, the notification provided by the software

application may include information to provide to the user. The information may, for instance, allow the user to make a more informed decision when prompted for confirmation. For instance, the user device may provide the natural-language query “I can get an uberX to your location in 9 minutes. Shall I request it?”

[0282] Thereafter, the user device causes (e.g., instructs) the software application to perform the task corresponding to the intent in accordance with the parameters.

[0283] At block 835, the user device receives a result response from the software application indicating whether the software application performed the task successfully. A result response indicating failure to perform a task may further indicate one or more reasons for the failure. In some examples, the user device may provide an output, such as a natural-language output, to the user indicating the one or more reasons for the failure.

[0284] A result response indicating success in performing a task may include one or more response items. Each response item may be a result determined (e.g., received, generated) by the software application when performing the task. By way of example, response items corresponding to reserving a car with a ride-booking application may include car type, license plate number, driver name, arrival time, current car location, pick-up location, destination, estimated trip time, estimated cost, and estimated trip route, service type (e.g., Uber pool vs. uberX). As another example, response items corresponding to initiating a workout session with a fitness application may include confirmation that the session has been initiated, workout duration, type of activity, and one or more goals.

[0285] In some examples, one or more of the response items may be provided to a user. With reference to FIG. 10B, for example, one or more of the response items may be provided to the user as a natural-language output 1012 as a text input and/or an audio output. Response items may be provided visually as well. For example, a map 1014 of an estimated trip route may be provided to the user. It will be appreciated that response items may be provided to users in any desired manner.

[0286] In some examples, the software application may specify the manner in which one or more response items are provided to the user. That is, the software application may determine the manner in which response items are displayed and/or spoken to the user and the digital assistant may provide each of the response items accordingly.

[0287] In some examples, the software application may specify how response items are provided using a UI extension of the digital assistant. The user device may, for instance, provide the software application with a set of view controller parameters (e.g., fields that can be supplied to a view controller for display), and in response, the software application may provide a set of view controller parameter values. The set of view controller parameter values may indicate which response items are to be displayed in various fields of a view controller and/or the manner in which the response items are displayed in each field.

[0288] In other examples, the digital assistant may determine how response items are provided. In yet other examples, the software application is invoked such that a user may interact with the software application directly. Invoking the application in this manner may terminate a session with the digital assistant in some examples.

[0289] In some examples, permissions of a software application are verified, for instance, prior to an intent and parameters being provided to the software application. The user device may, for instance, determine whether a software application is permitted to access data associated with a particular intent. The determination may be made based on permissions configured on the user device. In the example “Get me a black car at my location,” “location” may be a contextual parameter requiring contextual information (e.g., location data) of the user device. Thus, prior to resolving the location of the user device and providing the location as a parameter to the software application, the user device may first determine whether the software application is permitted to access the information. If the software application is permitted to access the data, operation proceeds as described. If the software application is not permitted to access the data, the intent and parameters are not provided to the software application and the task is not performed.

[0290] In some examples, a natural-language input may include a plurality of task requests. Thus, based on the natural-language input, a plurality of intents and/or a plurality of applications may be identified. Parameters associated with each of the intents are optionally identified as well. The natural-language input “Get me a ride to the airport and tell me my flight status,” for instance, may include both an intent to reserve a car and an intent to retrieve status of a user’s flight. In some examples, tasks corresponding to each of the intents may be performed sequentially or concurrently.

[0291] In some examples, a natural-language input may include a plurality of related task requests. For instance, in some examples, a requested task of a natural-language user input may depend on completion of another requested task of the natural-language input. In the example “Email me directions to the airport,” two tasks are requested: a first task to provide directions and a second task to email. The parameter “me” specifies a particular contact and is a parameter for the intent to email, and the parameter “airport” is a parameter specifying a destination for the intent to provide directions. Because emailing the directions requires that the directions first be provided, the second task (email) depends on the first task (provide directions). Accordingly, the task to retrieve directions is performed first.

[0292] In some examples, intents may be provided between applications. For example, an application may provide an intent object to another application to cause the application to perform a task. In the instant example, both intents (provide directions and email) may be provided to a maps application to provide the requested directions. The second intent, or the intent to email, may be provided to the maps application as a parameter, for instance. In accordance with the first intent, or the intent to provide directions, the maps application may provide the requested directions. Subsequently, the maps application may provide the second intent to the email application, for instance as an intent object including the directions as parameters. In response, the email application may email the directions as requested.

[0293] As another example, a user may provide a user input “Get me a ride to the Sharks game” while viewing a sports application (e.g., ESPN application). In response, the sports application, having information pertaining to the game, may pass an intent (e.g., reserve car) and a parameter

(e.g., address of game) to a ride-booking application. The intent and parameter may be provided as an intent object in some examples.

[0294] As yet another example, a user may provide a user input “Pay my brother \$5.00,” while using a ride-booking application. In response, the ride-booking application may pass an intent (e.g., pay) and a parameter (\$5.00) to a payment application (e.g., PayPal, Venmo, online payment service). As described, the intent and parameter may be provided as an intent object in some examples.

[0295] FIG. 9 illustrates a flow diagram of a process for operating a digital assistant in accordance with some embodiments. The process 900 may, for instance, be used to implement at least a portion of process 800 of FIG. 8, including, but not limited to, block 815 and/or block 820 of FIG. 8. Process 900 is performed, for example, using one or more electronic devices (e.g., devices 104, 108, 200, 400, or 600) implementing a digital assistant. In some examples, the process 900 is performed using a client-server system (e.g., system 100), and the blocks of the process 900 can be divided up in any manner between the server (e.g., DA server 106) and a client device. Thus, while portions of the process 900 are described herein as being performed by particular devices of a client-server system, it will be appreciated that the process 900 is not so limited. In other examples, the process 900 is performed using only a client device (e.g., user device 104). In process 900, some blocks are, optionally, combined, the order of some blocks is, optionally, changed, and some blocks are, optionally, omitted. In some examples, additional steps may be performed in combination with the process 900.

[0296] At block 905, a natural-language user input is received by a user device, such as the user device 104 of FIG. 1. As described, the natural-language input may include a request for the user device, and/or another device, to perform a task, and may further specify one or more parameters for the requested task.

[0297] At block 910, an intent and, optionally, one or more parameters associated with the intent, are identified. The intent and parameters may be derived from the natural-language user input. As previously described, intents may correspond to any type of tasks performed by a user device and, in particular, may correspond to tasks performed by one or more applications of a user device. Parameters associated with the intent may identify portions of the natural-language input that specify a manner in which a task corresponding to the intent is to be performed. In the example, “Get me a ride to the airport,” the intent corresponds to a task to reserve a car and “airport” is a parameter specifying a destination. Given that the intent of the user is to reserve a car, the location of the user may be a parameter (e.g., inferred parameter) as well.

[0298] At block 915, a determination is made as to whether a task corresponding to the intent can be fulfilled. In some examples, the determination may include determining whether an application configured to perform the task in accordance with the parameters is accessible to the user device. In the instant example, a determination is made as to whether the user device can access an application configured to reserve a car at a location of the user device. As described, accessible applications are those locally stored on the user device as well as those remotely accessible by the user device.

[0299] In accordance with a determination that the task corresponding to the intent can be fulfilled, at block 920, the intent and the parameter are provided to a software application. For example, if, at block 915, it is determined that the user device can access a software application configured to perform the requested task according to any identified parameters, the intent and parameters are provided to the application to perform the task. In the instant example, this includes determining that a software application configured to reserve a car at a location of the user device is accessible to the user device. For instance, the ride-booking application Lyft may be installed on the user device and may be used to reserve a car in accordance with examples described herein.

[0300] In accordance with a determination that the task corresponding to the intent cannot be fulfilled, at block 925, a list of one or more software applications is provided. The list of one or more software applications may, for instance, include one or more software applications configured to perform the task associated with the intent according to any identified parameters. In some examples, one or more software application of the list may be identified, for instance, based on one or more domains associated with the intent (recall that applications may be registered with one or more domains). With reference to FIG. 10C, once the list of software applications (e.g., ride-booking applications) has been determined, the user device may provide the list to the user. As illustrated, providing the list may include providing a natural-language input 1022 requesting the user select an application from the list of one or more software applications. In some examples, the list is generated by the user device. In other examples, the list of software applications is generated by a server and provided to the user device, which may in turn provide the list to a user, as described.

[0301] At block 930, the user device receives a user input indicative of a selection of one or more software applications of the list of one or more software applications. The user input may be a touch input on a touch-sensitive display of the user device and/or may be a natural-language user input.

[0302] At block 935, the intent and parameters are provided to the software application selected by the user. In some examples, providing the intent and parameters includes downloading and/or installing the software application such that the user device may access the software application locally. In other examples, this includes remotely accessing the selected application.

[0303] As described, in response to the intent and parameters, the software application may provide one or more responses. Once parameters have been validated, the user device may confirm intent with the software application and cause the software application to perform the task. The user device may thereafter receive result responses and, optionally, provide one or more response items of the result response to the user.

[0304] Reference is made herein to providing a user of a user device with natural-language outputs and/or natural language queries. In some examples, the manner in which natural-language outputs and queries are provided to a user may depend on a type or state of the user device. If, for instance, the user device is a mobile phone, the user device may provide a query using both text and audio. If, on the other hand, the user device is a speaker, the user device may provide the query using only audio. As another example, if the user device is a mobile phone not paired with head-

phones, the user device may provide a query using text and/or a relatively short natural-language query. If the user device is paired with headphones, the user device may provide the query using only a relatively long natural-language query.

[0305] FIGS. 10D-10F illustrate exemplary data flows of digital assistant systems in accordance with some embodiments. Data flows of FIGS. 10D-10F may be implemented using one or more of processes 800, 900 in some examples. FIG. 10D illustrates an exemplary data flow of a digital assistant system 1030 in accordance with some embodiments. In particular, FIG. 10D illustrates a data flow for an application registration process and a data flow for performance of a task. Data flows 1031-1041 may be associated with the application registration process (illustrated in FIG. 10D as hashed lines) and data flows 1043-1051 may be associated with performance of a task (illustrated in FIG. 10D as solid lines).

[0306] Generally, data flows associated with the application registration process are directed to registration of an application with an application registration service (e.g., validation service) by which applications, and custom vocabulary corresponding thereto, can be accessed and/or employed by a digital assistant to perform tasks.

[0307] In operation, at data flow 1031, an application is submitted to an application review module 1032. Both a language model corresponding to the application and intents of the application may be submitted as well. The language model may include custom vocabulary for the application. In turn, at data flow 1033, the application review module 1032 may provide the application, custom vocabulary, and/or intents of the application to the validation service 1034. The validation service 1034 may determine whether to validate the application, for instance, based on whether the application is operable with a digital assistant. This may include for instance, ensuring that any intents of the application correspond to one or more domains of the application. For instance, the validation service may reject a messaging application associated with intents to reserve cars because the domain and intents mismatch. The validation service 1034, at data flow 1035, may provide a validation response indicating whether the application is valid.

[0308] If the validation service 1034 indicates that the application is valid, the application review module 1032 provides the application (as verified) to the application store 1036. Generally, the application may be downloaded and/or accessed at the application store 1036 by the user device 1040 via the DA server 1038, as indicated by the data flow 1039. The user device 1040 may be the user device 104 of FIG. 1 and the DA server 1038 may be the DA server 106 of FIG. 1 in some examples. This may, for instance, result in an application list (e.g., info.plist) of the user device 1040 being updated and/or synchronized with the DA server 1038. At data flow 1041, the validation service may provide custom vocabulary (e.g., runtime vocabulary) of the application to the DA server 1038 to assist in parsing natural-language inputs, as described.

[0309] Generally, data flows associated with performance of a task are directed to providing an intent and, optionally, one or more parameters to an application for performing a task corresponding to the intent.

[0310] In operation, at data flow 1043, the user device 1040 may provide a natural-language input to the DA server 1038. The natural-language input may be provided by the

digital assistant 1042 of the user device 1040 in some examples. Based on the natural-language input, the DA server 1038 may identify one or more tasks requested in the natural-language input and one or more parameters associated with the intent. Additionally, the DA server 1038 may identify an application for performing a task with the intent. A name (or other form of identifier) of the identified application 1044 may be a parameter of the intent in some examples. At data flow 1045, the DA server 1038 thereafter provides the intent, parameters, and an identification of the identified application 1044 to the user device 1040 (e.g., the digital assistant 1042 of the user device 1040). The intent and parameters may be provided to the user device 1040 as an intent object in some examples.

[0311] In response, the digital assistant determines whether the identified application 1044 is permitted to access information associated with the identified parameters. For instance, if a parameter is a location of the user device 1040, the digital assistant queries data permissions 1046 to determine if the application 1044 is permitted to access location data.

[0312] In the event that the application is permitted to access data for each of the parameters, the user device (e.g., the digital assistant 1042 of the user device) provides the intent to the application. As illustrated, the application may reside on the user device 1040. In other examples, the application may reside on one or more other devices, and the intent may be transmitted to the application over one or more networks. As described, the application 1044 may thereafter request input from a user of the user device 1040 if the application determines that one or more parameters are missing, improper, and/or ambiguous. In some examples, queries for user input may be provided as natural-language queries generated by the DA server 1038. Accordingly, at data flow 1051, the user device 1040 may request, and thereafter, receive, one or more natural-language queries. Once all parameters have been resolved, the application 1044 may perform a task corresponding to the intent and provide a result response indicating whether the task was successfully performed.

[0313] One or more of the data flows of FIG. 10D are implemented (e.g., generated), for example, using one or more electronic devices (e.g., devices 104, 108, 200, 400, or 600) implementing a digital assistant. In particular, data flows provided between DA server 1038 and the digital assistant 1042 of the user device 1040 are illustrated as relying on a client-server architecture. In other examples, the DA server 1038 may be implemented as a process and/or service on the user device 1040. Accordingly, data flows exchanged between the DA server 1038 and the digital assistant 1042 may be exchanged solely on the user device 1040 in some examples.

[0314] FIG. 10E illustrates an exemplary data flow of a digital assistant system 1060 in accordance with some embodiments. In particular, FIG. 10E illustrates an exemplary data flow for an application registration process and may be used to implement the application registration process as set forth in FIG. 10D. Further, several elements of FIG. 10E correspond to elements of FIG. 10D, respectively, and have been provided with the same reference numerals. In the interest of brevity, an explanation of their function and operation will not be repeated.

[0315] At data flow 1065, validated vocabulary is provided from validation service 1034 to global application

vocabulary store **1060**. Generally, the global application vocabulary store may store language models and/or vocabulary for any number and/or versions of software applications. At data flows **1061** and **1063**, speech training module **1062** and natural-language training module **1064** are trained to recognize and natural language process the application specific vocabulary supplied with the validated application. Based on the data, the global application vocabulary store may generate and/or train one or more language models that allow the digital assistant to recognize and process utterances containing the application specific vocabulary.

[0316] During operation of a user device, such as the user device **1040** of FIG. 10D, the runtime-specific global application vocabulary store may receive vocabulary and/or language models for one or more applications of the user device from the global application vocabulary store **1060**. The vocabulary may be specific to a user ID of a user of the user device, and/or may be specific to versions of applications and/or operating systems of the user device. Based on the vocabulary, one or more terms of natural-language inputs may be identified, for instance, as parameters.

[0317] In accordance with some embodiments, FIG. 11 shows a functional block diagram of an electronic device **1100** configured in accordance with the principles of the various described embodiments, including those described with reference to FIG. 8. The functional blocks of the device are, optionally, implemented by hardware, software, or a combination of hardware and software to carry out the principles of the various described embodiments. It is understood by persons of skill in the art that the functional blocks described in FIG. 11 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein.

[0318] As shown in FIG. 11, an electronic device **1100** includes a touch-sensitive display unit **1102** and a processing unit **1108** optionally coupled to touch-sensitive display unit **1102**. In some embodiments, the processing unit **1108** includes a receiving unit **1110**, an identifying unit **1112**, a providing unit **1114**, and optionally, an outputting unit **1116**, a speaking unit **1118**, a display enabling unit **1120**, a determining unit **1122**, a requesting unit **1124**, an accessing unit **1126**, and a causing unit **1128**.

[0319] In some examples, the processing unit **1108** is configured to receive (e.g., with the receiving unit **1110**) a natural-language user input (e.g., block **805** of FIG. 8); identify (e.g., with identifying unit **1112**), with the one or more processors, an intent object of a set of intent objects and a parameter associated with the intent object, where the intent object and the parameter are derived from the natural-language user input (e.g., block **810** of FIG. 8); identify (e.g., with identifying unit **1112**) a software application associated with the intent object of the set of intent objects (e.g., block **815** of FIG. 8); and provide (e.g., with the providing unit **1114**) the intent object and the parameter to the software application (e.g., block **820** of FIG. 8).

[0320] In some examples, the processing unit **1108** is further configured to receive (e.g., with the receiving unit **1110**) a response from the software application, wherein the response is indicative of whether the parameter is valid (e.g., block **825** of FIG. 8).

[0321] In some examples, receiving (e.g., with the receiving unit **1110**) a response from the software application comprises receiving (e.g., with the receiving unit **1110**), from the software application, a disambiguation list associated with the parameter, where the disambiguation list includes a plurality of candidate parameters.

[0322] In some examples, the processing unit **1108** is further configured to output (e.g., with the outputting unit **1116**) the disambiguation list; receive (e.g., with the receiving unit **1110**) a user input indicative of a selection of a candidate parameter value of the disambiguation list; and provide (e.g., with the providing unit **1114**) the selected candidate parameter value of the disambiguation list to the software application.

[0323] In some examples, the natural-language user input is a first natural-language input, the parameter is a first parameter, and the processing unit **1108** is further configured to receive (e.g., with the receiving unit **1110**), from the software application, a request for a second parameter associated with the intent object; provide (e.g., with the providing unit **1114**) a natural-language query based on the request; receive (e.g., with the receiving unit **1110**) a second natural-language user input; identify (e.g., with identifying unit **1112**) the second parameter, where the second parameter is derived from the second natural-language user input; and provide (e.g., with the providing unit **1114**) the second parameter to the software application.

[0324] In some examples, the electronic device further comprises an audio output component, and providing the natural-language query comprises speaking (e.g., with the speaking unit **1118**), via speech synthesis, the natural-language query through the audio output component.

[0325] In some examples, providing the natural-language query comprises providing (e.g., with the providing unit **1114**) the natural-language query based on a type of the electronic device, a state of the electronic device, or a combination thereof.

[0326] In some examples, the processing unit **1108** is further configured to, after providing the intent object and the parameter to the software application, receive (e.g., with the receiving unit **1110**), from the software application, a result response associated with the intent object (e.g., block **835** of FIG. 8).

[0327] In some examples, receiving a result response comprises receiving (e.g., with the receiving unit **1110**), from the software application, a set of response items associated with the intent object, and outputting (e.g., with the outputting unit **1116**) the set of response items.

[0328] In some examples, the processing unit **1108** is further configured to provide (e.g., with the providing unit **1114**), to the software application, a set of view controller parameters; and receive (e.g., with the receiving unit **1110**), from the software application, a set of view controller parameter values corresponding to the set of view controller parameters, where outputting the set of response items comprises enabling display (e.g., with the display enabling unit **1120**) of the set of response items in a user interface based on the received view controller parameter values.

[0329] In some examples, each intent object of the set of intent objects is associated with the software application.

[0330] In some examples, the intent object is a first intent object and wherein the parameter is associated with the first intent object and not associated with a second intent object.

[0331] In some examples, the processing unit **1108** is further configured to determine (e.g., with the determining unit **1122**) whether the electronic device is in a lock state; and in accordance with a determination that the device is in the lock state, determine (e.g., with the determining unit **1122**) whether the intent object is allowed to be provided to the software application while the electronic device is in the lock state, where the intent object and the parameter value are provided to the software application only in accordance with a determination that the intent object is allowed to be provided to the software application while the electronic device is in the lock state.

[0332] In some examples, the parameter is derived from the natural-language user input by analyzing the natural-language user input using each of a plurality of detectors.

[0333] In some examples, the processing unit **1108** is further configured to determine (e.g., with the determining unit **1122**) a user context of the electronic device, where the parameter is based at least in part on the user context.

[0334] In some examples, the processing unit **1108** is further configured to request (e.g., with the requesting unit **1124**) a confirmation of the parameter; and receive (e.g., with the receiving unit **1110**) a user input corresponding to the confirmation of the parameter, where the parameter is provided to the software application in response to receiving the user input.

[0335] In some examples, identifying the software application comprises determining (e.g., with the determining unit **1122**) whether the software application resides on the electronic device; and in accordance with a determination that the software application does not reside on the electronic device, determining (e.g., with the determining unit **1122**) whether the software application resides on an external device that is in communication with the electronic device. In some examples, providing the intent object and the parameter to the software application comprises, in accordance with a determination that the software application resides on an external device that is in communication with the electronic device, providing (e.g., with the providing unit **1114**) the intent object and the parameter to the software application by providing the intent object and the parameter to the external device.

[0336] In some examples, identifying the software application comprises determining (e.g., with the determining unit **1122**) whether the software application resides on the electronic device; in accordance with a determination that the software application does not reside on the electronic device, identifying (e.g., with the identifying unit **1112**) a set of software applications associated with the intent object; enabling display (e.g., with the display enabling unit **1120**) of, in a user interface, the set of software applications associated with the intent object; receiving (e.g., with the receiving unit **1110**) a user input including a selection of one or more software applications of the set of software applications associated with the intent object; and accessing (e.g., with the accessing unit **1126**) the selected one or more software applications of the set of software applications associated with the intent object.

[0337] In some examples, identifying a software application associated with the intent object of the set of intent objects comprises identifying (e.g., with the identifying unit **1112**) the software application based on the parameter.

[0338] In some examples, the processing unit **1108** is further configured to identify (e.g., with the identifying unit

1112), based on the natural-language user input, a third intent object from the set of intent objects; identify (e.g., with the identifying unit **1112**), based on the identified third intent object, a second software application; provide (e.g., with the providing unit **1114**) the third intent object and at least one response item to the second software application; and receive (e.g., with the receiving unit **1110**), from the second software application, a second response item associated with the third intent object.

[0339] In some examples, the software application is a first software application, and providing the intent object and the parameter to the software application comprises causing (e.g., with the causing unit **1128**) a third software application to provide the intent object and the parameter to the first software application.

[0340] In some examples, receiving a natural-language user input comprises receiving (e.g., with the receiving unit **1110**) a natural-language user input including an application-specific term; and where identifying a software application associated with the intent object of the set of intent objects comprises identifying (e.g., with the identifying unit **1112**) the software application based on the application-specific term.

[0341] In some examples, the processing unit **1108** is further configured to receive (e.g., with the receiving unit **1110**) a command including an identification of the software application; and in response to the command, determine (e.g., with the determining unit **1122**) whether the software application is permitted to access data associated with the intent object.

[0342] In some examples, the processing unit **1108** is further configured to cause (e.g., with the causing unit **1128**) the software application to perform a task associated with the intent object based on the parameter.

[0343] The operation described above with respect to FIG. 8 is, optionally, implemented by components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 11. For example, receiving operations **805**, **825**, and **835**; identifying operations **810** and **815**, providing operation **820**, and confirming operation **830** are optionally implemented by processor(s) **120**. It would be clear to a person of ordinary skill in the art how other processes can be implemented based on the components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 11.

[0344] It is understood by persons of skill in the art that the functional blocks described in FIG. 11 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein. For example, processing unit **1108** can have an associated “controller” unit that is operatively coupled with processing unit **1108** to enable operation. This controller unit is not separately illustrated in FIG. 11 but is understood to be within the grasp of one of ordinary skill in the art who is designing a device having a processing unit **1108**, such as device **1100**. As another example, one or more units, such as the receiving unit **1110**, may be hardware units outside of processing unit **1108** in some embodiments. The description herein thus optionally supports combination, separation, and/or further definition of the functional blocks described herein.

[0345] In accordance with some embodiments, FIG. 12 shows a functional block diagram of an electronic device **1200** configured in accordance with the principles of the

various described embodiments, including those described with reference to FIG. 8. The functional blocks of the device are, optionally, implemented by hardware, software, or a combination of hardware and software to carry out the principles of the various described embodiments. It is understood by persons of skill in the art that the functional blocks described in FIG. 12 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein.

[0346] As shown in FIG. 12, one or more electronic devices 1200 include one or more processing units 1208. In some embodiments, the one or more processing units 1208 include a receiving unit 1210, a determining unit 1212, an identifying unit 1214, and a providing unit 1216.

[0347] In some examples, the one or more processing units 1208 is configured to receive (e.g., with the receiving unit 1210) a natural-language user input (e.g., block 805 of FIG. 8); determine (e.g., with the determining unit 1212), based on the natural-language user input, an intent object of a set of intent objects and a parameter associated with the intent object (e.g., block 810 of FIG. 8); identify (e.g., with the identifying unit 1214) a software application based on at least one of the intent object or the parameter (e.g., block 815 of FIG. 8); and provide (e.g., with the providing unit 1216) the intent object and the parameter to the software application (e.g., block 820 of FIG. 8).

[0348] In some examples, determining, based on the natural-language user input, an intent object of a set of intent objects and a parameter associated with the intent object comprises determining (e.g., with the determining unit 1212), with a first electronic device of the one or more electronic devices 1200, the intent object of the set of intent objects and the parameter associated with the intent object; and where providing the intent object and the parameter to the software application comprises providing (e.g., with the providing unit 1216), with a second electronic device of the one or more electronic devices 1200, the intent object and the parameter to the software application.

[0349] In some examples, the one or more processing units 1208 are configured to provide (e.g., with the providing unit 1216), with the first electronic device of the one or more electronic devices 1200, a command to the second electronic device of the one or more electronic devices 1200; and in response to the command, determine (e.g., with the determining unit 1212), with the second electronic device, whether the software application is permitted to access data associated with the intent object.

[0350] In some examples, the one or more processing units 1208 are configured to receive (e.g., with the receiving unit 1210) a response from the software application, where the response is indicative of whether the parameter is valid (e.g., block 825 of FIG. 8).

[0351] In some examples, the parameter is indicative of the software application.

[0352] In some examples, the one or more processing units 1208 are configured to identify (e.g., with the identifying unit 1214) the software application based on the natural-language user input; and determine (e.g., with the determining unit 1212) whether the intent object of the set of intent objects corresponds to a registered intent object of the software application; where the intent object and the

parameter are determined only in accordance with a determination that the intent object corresponds to a registered intent object of the software application.

[0353] In some examples, the parameter is a first parameter, and the one or more processing units are configured to determine (e.g., with the determining unit 1212), based on the natural-language user input, a second parameter; and provide (e.g., with the providing unit 1216) the second parameter to the software application.

[0354] In some examples, the response is a first response, and the one or more processing units are configured to receive (e.g., with the receiving unit 1210) a second response from the software application, where the second response is indicative of whether the second parameter is valid.

[0355] The operation described above with respect to FIG. 8 is, optionally, implemented by components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 12. For example, receiving operations 805, 825, and 835; identifying operations 810 and 815, providing operation 820, and confirming operation 830 are optionally implemented by processor(s) 120. It would be clear to a person of ordinary skill in the art how other processes can be implemented based on the components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 12.

[0356] It is understood by persons of skill in the art that the functional blocks described in FIG. 12 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein. For example, the one or more processing units 1208 can have an associated “controller” unit that is operatively coupled with at least one of the one or more processing units 1208 to enable operation. This controller unit is not separately illustrated in FIG. 12 but is understood to be within the grasp of one of ordinary skill in the art who is designing a device having one or more processing units 1208, such as device 1200. As another example, one or more units, such as the receiving unit 1210, may be hardware units outside of the one or more processing units 1208 in some embodiments. The description herein thus optionally supports combination, separation, and/or further definition of the functional blocks described herein.

[0357] In accordance with some embodiments, FIG. 13 shows a functional block diagram of an electronic device 1300 configured in accordance with the principles of the various described embodiments, including those described with reference to FIG. 9. The functional blocks of the device are, optionally, implemented by hardware, software, or a combination of hardware and software to carry out the principles of the various described embodiments. It is understood by persons of skill in the art that the functional blocks described in FIG. 13 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein.

[0358] As shown in FIG. 13, one or more electronic devices 1300 include one or more processing units 1308. In some embodiments, the one or more processing units 1308 include a receiving unit 1310, an identifying unit 1312, a determining unit 1314, a providing unit 1316, and optionally, a causing unit 1318.

[0359] In some examples, the one or more processing units **1308** are configured to receive (e.g., with the receiving unit **1310**) a natural-language user input (e.g., block **905** of FIG. 9); identify (e.g., with the identifying unit **1312**), based on the natural-language user input, an intent object of a set of intent objects and a parameter associated with the intent object (e.g., block **910** of FIG. 9); determine (e.g., with the determining unit **1314**), based on at least one of the intent object or the parameter, whether a task corresponding to the intent object can be fulfilled (e.g., block **915** of FIG. 9); in accordance with a determination that the task corresponding to the intent object can be fulfilled, provide (e.g., with the providing unit **1316**) the intent object and the parameter to a software application associated with the intent object (e.g., block **920** of FIG. 9); and in accordance with a determination that the task corresponding to the intent object cannot be fulfilled, provide (e.g., with the providing unit **1316**) a list of one or more software applications associated with the intent object (e.g., block **925** of FIG. 9).

[0360] In some examples, the one or more processing units **1308** are configured to, after providing a list of one or more software applications associated with the intent object, receive (e.g., with the receiving unit **1310**) a user input indicative of a selection of a software application of the list of one or more software applications (e.g., block **930** of FIG. 9); and in response to the user input, provide (e.g., with the providing unit **1316**) the intent object of the set of intent objects to the selected software application (e.g., block **935** of FIG. 9).

[0361] In some examples, the one or more processing units **1308** are configured to, in response to the user input, provide (e.g., with the providing unit **1316**) the parameter to the selected software application.

[0362] In some examples, the one or more processing units **1308** are configured to receive (e.g., with the receiving unit **1310**) a response from the selected software application, and the response is indicative of whether the parameter is valid.

[0363] In some examples, the one or more processing units **1308** are configured to cause (e.g., with the causing unit **1318**) the selected software application to perform the task corresponding to the intent object, and after providing the intent object to the selected software application, receive (e.g., with the receiving unit **1310**), from the selected software application, a result response associated with the intent object.

[0364] The operation described above with respect to FIG. 9 is, optionally, implemented by components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 13. For example, determining operation **905**, providing operations **910**, **915**, and **925**; and receiving operation **920**, are optionally implemented by processor(s) **120**. It would be clear to a person of ordinary skill in the art how other processes can be implemented based on the components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 13.

[0365] It is understood by persons of skill in the art that the functional blocks described in FIG. 13 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein. For example, the one or more processing units **1308** can have an associated “controller” unit that is operatively coupled with at least one of the one

or more processing units **1308** to enable operation. This controller unit is not separately illustrated in FIG. 13 but is understood to be within the grasp of one of ordinary skill in the art who is designing a device having one or more processing units **1308**, such as device **1300**. As another example, one or more units, such as the receiving unit **1310**, may be hardware units outside of the one or more processing units **1308** in some embodiments. The description herein thus optionally supports combination, separation, and/or further definition of the functional blocks described herein.

[0366] In accordance with some embodiments, FIG. 14 shows a functional block diagram of an electronic device **1400** configured in accordance with the principles of the various described embodiments, including those described with reference to FIG. 9. The functional blocks of the device are, optionally, implemented by hardware, software, or a combination of hardware and software to carry out the principles of the various described embodiments. It is understood by persons of skill in the art that the functional blocks described in FIG. 14 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein.

[0367] As shown in FIG. 14, an electronic device **1400** includes a touch-sensitive display unit **1402** and a processing unit **1408** optionally coupled to touch-sensitive display unit **1402**. In some embodiments, the processing unit **1402** includes a receiving unit **1410**, a providing unit **1412**, an obtaining unit **1414**, a display enabling unit **1416**, and optionally, an identifying unit **1418** and a causing unit **1420**.

[0368] In some examples, the processing unit **1408** is configured to receive (e.g., with the receiving unit **1410**) a natural-language user input, where the natural-language user input is indicative of an intent object of a set of intent objects (e.g., block **905** of FIG. 9); provide (e.g., with the providing unit **1412**) the natural-language user input to a second electronic device; receive (e.g., with the receiving unit **1410**), from the second electronic device, an indication that a software application associated with the intent object is not located on the first electronic device (e.g., block **925** of FIG. 9); responsive to the notification, obtain (e.g., with the obtaining unit **1414**) the list of applications associated with the intent object; enable display (e.g., with the display enabling unit **1416**) of, with a touch-sensitive display of the first electronic device, the list of applications associated with the intent object in a user interface; receive (e.g., with the receiving unit **1410**) a user input indicative of a selection of an application of the list of applications (e.g., block **930** of FIG. 9); and provide (e.g., with the providing unit **1412**) the intent object of the set of intent objects to the application (e.g., block **935** of FIG. 9).

[0369] In some examples, the processing unit **1408** is further configured to identify (e.g., with the identifying unit **1418**) the intent object of the set of intent objects.

[0370] In some examples, the processing unit **1408** is further configured to identify (e.g., with the identifying unit **1418**) a parameter associated with the intent object of the set of intent objects; and provide (e.g., with the providing unit **1412**) the parameter to the application.

[0371] In some examples, the processing unit **1408** is further configured to receive (e.g., with the receiving unit

[1410] a response from the application, and the response is indicative of whether the parameter is valid.

[0372] In some examples, the processing unit **1408** is further configured to cause (e.g., with the causing unit **1420**) the application to perform a task associated with the intent object; and after providing the intent object to the application, receive (e.g., with the receiving unit **1410**), from the application, a result response associated with the intent object.

[0373] The operation described above with respect to FIG. 9 is, optionally, implemented by components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 14. For example, determining operation **905**, providing operations **910**, **915**, and **925**; and receiving operation **920** are optionally implemented by processor(s) **120**. It would be clear to a person of ordinary skill in the art how other processes can be implemented based on the components depicted in FIGS. 1, 2A, 4, 6A-6B, 7A, and 14.

[0374] It is understood by persons of skill in the art that the functional blocks described in FIG. 14 are, optionally, combined or separated into sub-blocks to implement the principles of the various described embodiments. Therefore, the description herein optionally supports any possible combination or separation or further definition of the functional blocks described herein. For example, processing unit **1408** can have an associated “controller” unit that is operatively coupled with processing unit **1408** to enable operation. This controller unit is not separately illustrated in FIG. 14 but is understood to be within the grasp of one of ordinary skill in the art who is designing a device having a processing unit **1408**, such as device **1400**. As another example, one or more units, such as the receiving unit **1410**, may be hardware units outside of processing unit **1408** in some embodiments. The description herein thus optionally supports combination, separation, and/or further definition of the functional blocks described herein.

[0375] The foregoing description, for purpose of explanation, has been described with reference to specific embodiments. However, the illustrative discussions above are not intended to be exhaustive or to limit the invention to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. The embodiments were chosen and described in order to best explain the principles of the techniques and their practical applications. Others skilled in the art are thereby enabled to best utilize the techniques and various embodiments with various modifications as are suited to the particular use contemplated.

[0376] Although the disclosure and examples have been fully described with reference to the accompanying drawings, it is to be noted that various changes and modifications will become apparent to those skilled in the art. Such changes and modifications are to be understood as being included within the scope of the disclosure and examples as defined by the claims.

What is claimed is:

1. An electronic device, comprising:

one or more processors;
a memory; and

one or more programs, wherein the one or more programs are stored in the memory and configured to be executed by the one or more processors, the one or more programs including instructions for:

obtaining, from a software application, an application vocabulary for the software application, wherein the

application vocabulary includes at least a vocabulary entry of a first type and a vocabulary entry of a second type;

registering the application vocabulary to a knowledge base for a digital assistant of the electronic device;

receiving a user input;

determining whether the user input corresponds to a first vocabulary entry of the application vocabulary; and

in accordance with a determination that at least a first portion of the user input matches the first vocabulary entry, causing the software application to perform a first action based on the first vocabulary entry.

2. The electronic device of claim 1, wherein the vocabulary entry of the first type is different from the second type.

3. The electronic device of claim 1, wherein the first vocabulary entry represents a first object handled by the software application.

4. The electronic device of claim 1, wherein the first vocabulary entry represents a first type handled by the software application.

5. The electronic device of claim 1, wherein the first vocabulary entry is associated with a first command for the first action in the knowledge base.

6. The electronic device of claim 1, wherein determining whether the user input corresponds to the first vocabulary entry is performed while the software application is not in an active state.

7. The electronic device of claim 1, wherein causing the software application to perform the first action based on the first vocabulary entry includes:

identifying the first action using second metadata associated with the first vocabulary entry in the knowledge base.

8. The electronic device of claim 1, wherein causing the software application to perform the first action based on the first vocabulary entry includes:

identifying the software application using third metadata associated with the first vocabulary entry in the knowledge base.

9. The electronic device of claim 1, wherein obtaining the application vocabulary includes:

in response to receiving a user input requesting to install the software application, obtaining at least a first portion of the application vocabulary.

10. The electronic device of claim 1, wherein the software application is installed on the electronic device.

11. The electronic device of claim 1, wherein obtaining the application vocabulary for the software application includes:

in response to launching the software application, obtaining at least a second portion of the application vocabulary.

12. The electronic device of claim 1, wherein obtaining the application vocabulary for the software application includes:

receiving a request from the software application; and
in response to receiving the request from the software application, obtaining at least a third portion of the application vocabulary, wherein the third portion of the application vocabulary includes the vocabulary entry of the second type.

13. The electronic device of claim **12**, wherein the request from the software application is received as an application programming interface (API) call.

14. The electronic device of claim **1**, wherein registering the application vocabulary to the knowledge base includes: for each respective vocabulary entry included in the application vocabulary, associating respective metadata with the respective vocabulary entry in the knowledge base.

15. The electronic device of claim **1**, wherein a global vocabulary database stores the application vocabulary.

16. The electronic device of claim **1**, wherein the application vocabulary corresponds to a user identification for a user of the electronic device.

17. The electronic device of claim **1**, wherein the application vocabulary corresponds to an operating system of the electronic device.

18. The electronic device of claim **1**, wherein the application vocabulary corresponds to a software version.

19. A non-transitory computer-readable storage medium storing one or more programs, the one or more programs configured to be executed by one or more processors of an electronic device, the one or more programs including instructions for:

obtaining, from a software application, an application vocabulary for the software application, wherein the application vocabulary includes at least a vocabulary entry of a first type and a vocabulary entry of a second type;

registering the application vocabulary to a knowledge base for a digital assistant of the electronic device; receiving a user input;

determining whether the user input corresponds to a first vocabulary entry of the application vocabulary; and in accordance with a determination that at least a first portion of the user input matches the first vocabulary entry, causing the software application to perform a first action based on the first vocabulary entry.

20. A computer-implemented method, comprising at an electronic device with one or more processors, one or more display generation components, and memory:

obtaining, from a software application, an application vocabulary for the software application, wherein the application vocabulary includes at least a vocabulary entry of a first type and a vocabulary entry of a second type;

registering the application vocabulary to a knowledge base for a digital assistant of the electronic device; receiving a user input;

determining whether the user input corresponds to a first vocabulary entry of the application vocabulary; and in accordance with a determination that at least a first portion of the user input matches the first vocabulary entry, causing the software application to perform a first action based on the first vocabulary entry.

* * * * *