

(19) **United States**
(12) **Patent Application Publication** (10) **Pub. No.: US 2025/0260417 A1**
Galvin (43) **Pub. Date: Aug. 14, 2025**

(54) **SYSTEM AND METHOD FOR
COMPRESSING AND RESTORING DATA
USING HIERARCHICAL AUTOENCODERS
AND A KNOWLEDGE-ENHANCED
CORRELATION NETWORK**

(71) Applicant: **AtomBeam Technologies Inc.**, Moraga,
CA (US)

(72) Inventor: **Brian Galvin**, Silverdale, WA (US)

(21) Appl. No.: **19/192,182**

(22) Filed: **Apr. 28, 2025**

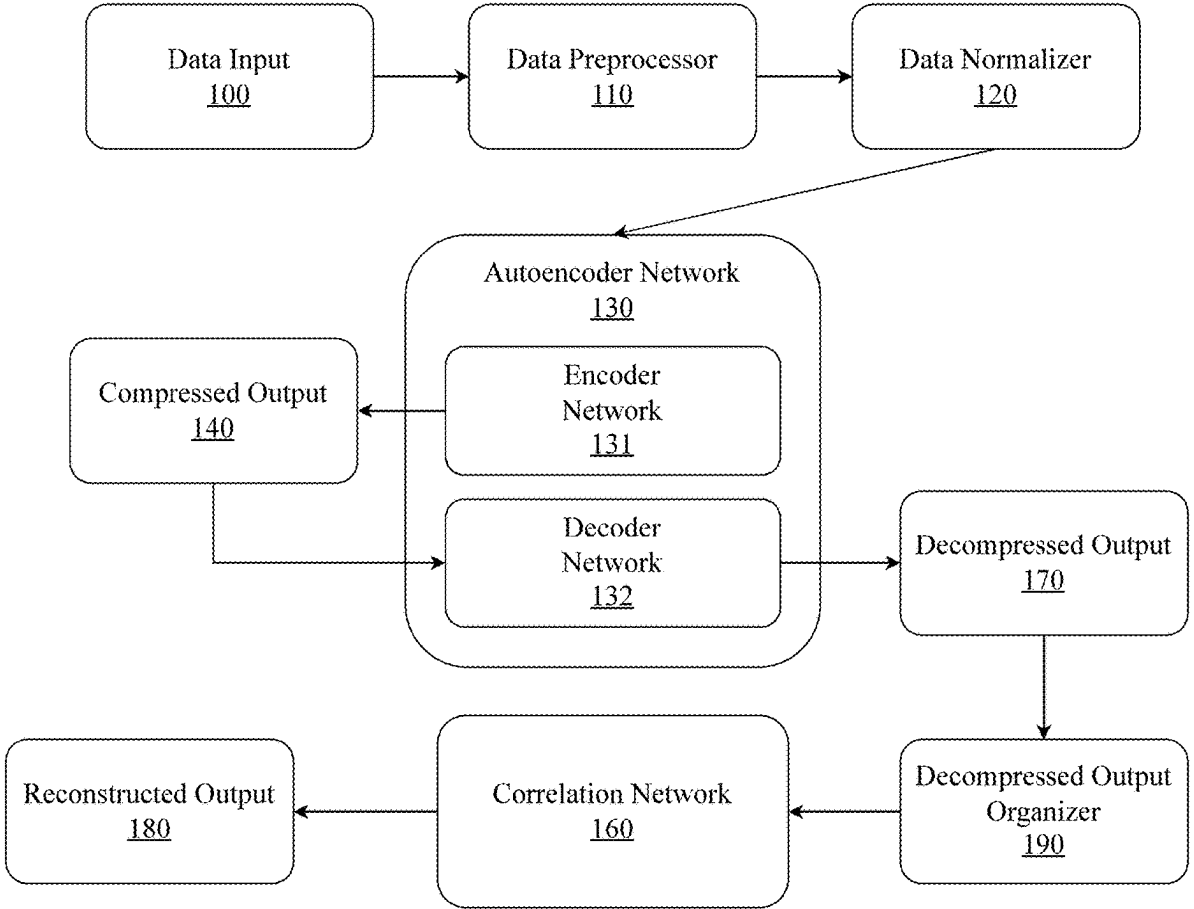
Publication Classification

(51) **Int. Cl.**
H03M 7/30 (2006.01)
G06N 3/0455 (2023.01)
G06N 3/08 (2023.01)

(52) **U.S. Cl.**
CPC **H03M 7/6011** (2013.01); **G06N 3/0455**
(2023.01); **G06N 3/08** (2013.01); **H03M**
7/6005 (2013.01); **H03M 7/70** (2013.01)

(21) **ABSTRACT**

A system and method for compressing and restoring data using a hierarchical multi-level autoencoder architecture and correlation network. The system compresses data using a cascade of autoencoders, each focusing on different scales or features, allowing for efficient representation across various resolutions. Data restoration employs a corresponding hierarchical decoder structure and a correlation network, trained on cross-correlated data sets. This approach leverages inter-data relationships at multiple scales, potentially recovering more lost information than traditional single-scale methods. The hierarchical structure adapts to diverse data types, achieving higher compression ratios while maintaining data quality, applicable to fields such as remote sensing, IoT data processing, and multimedia compression.



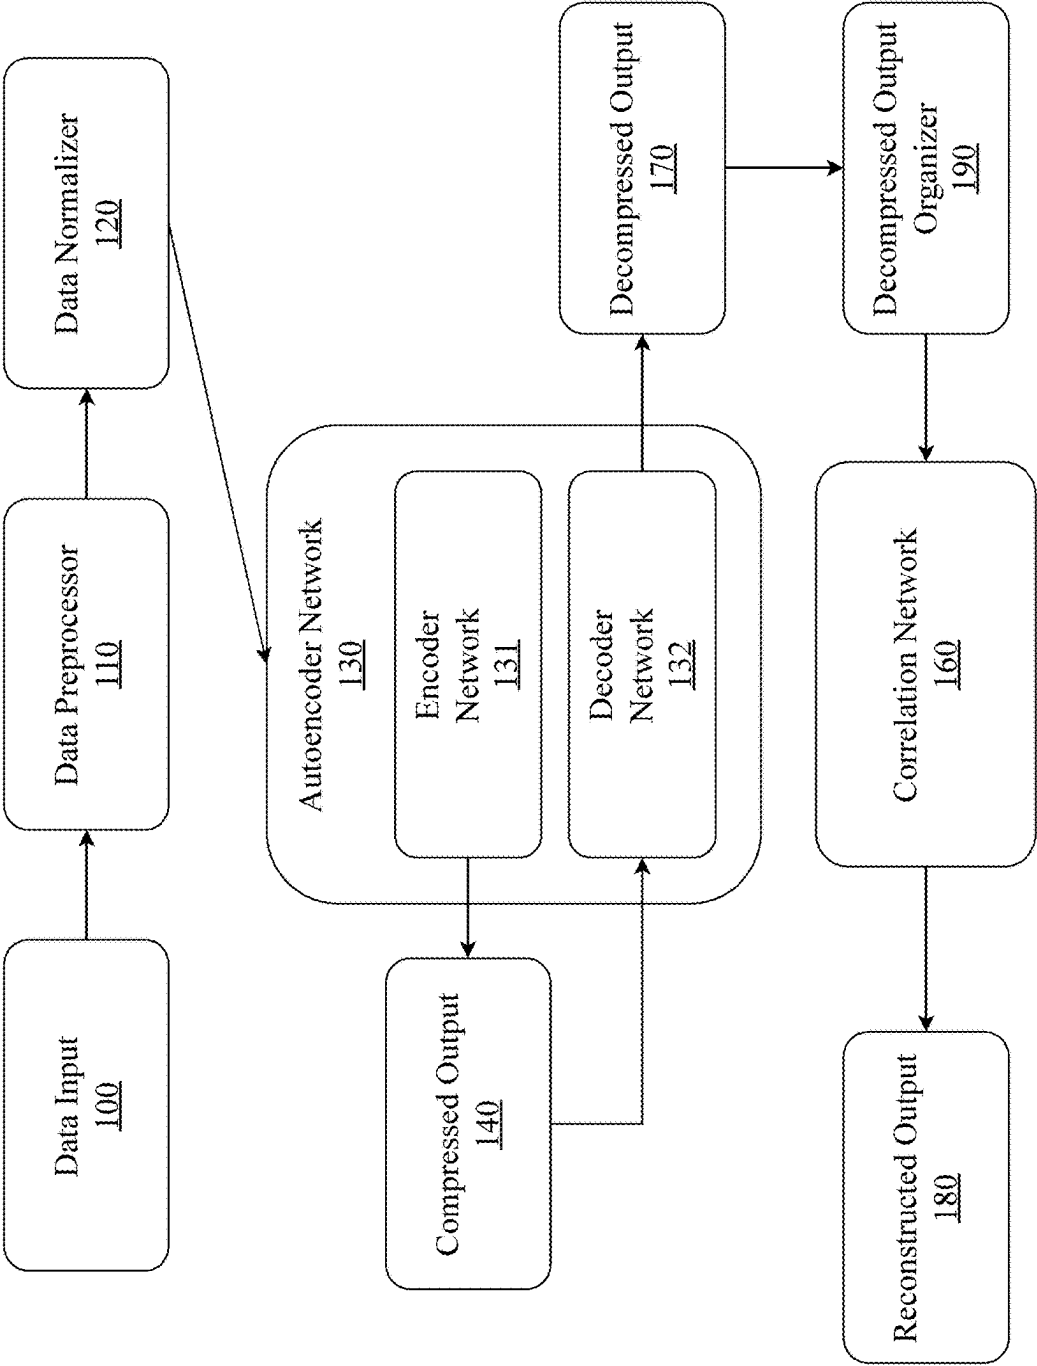


FIG. 1

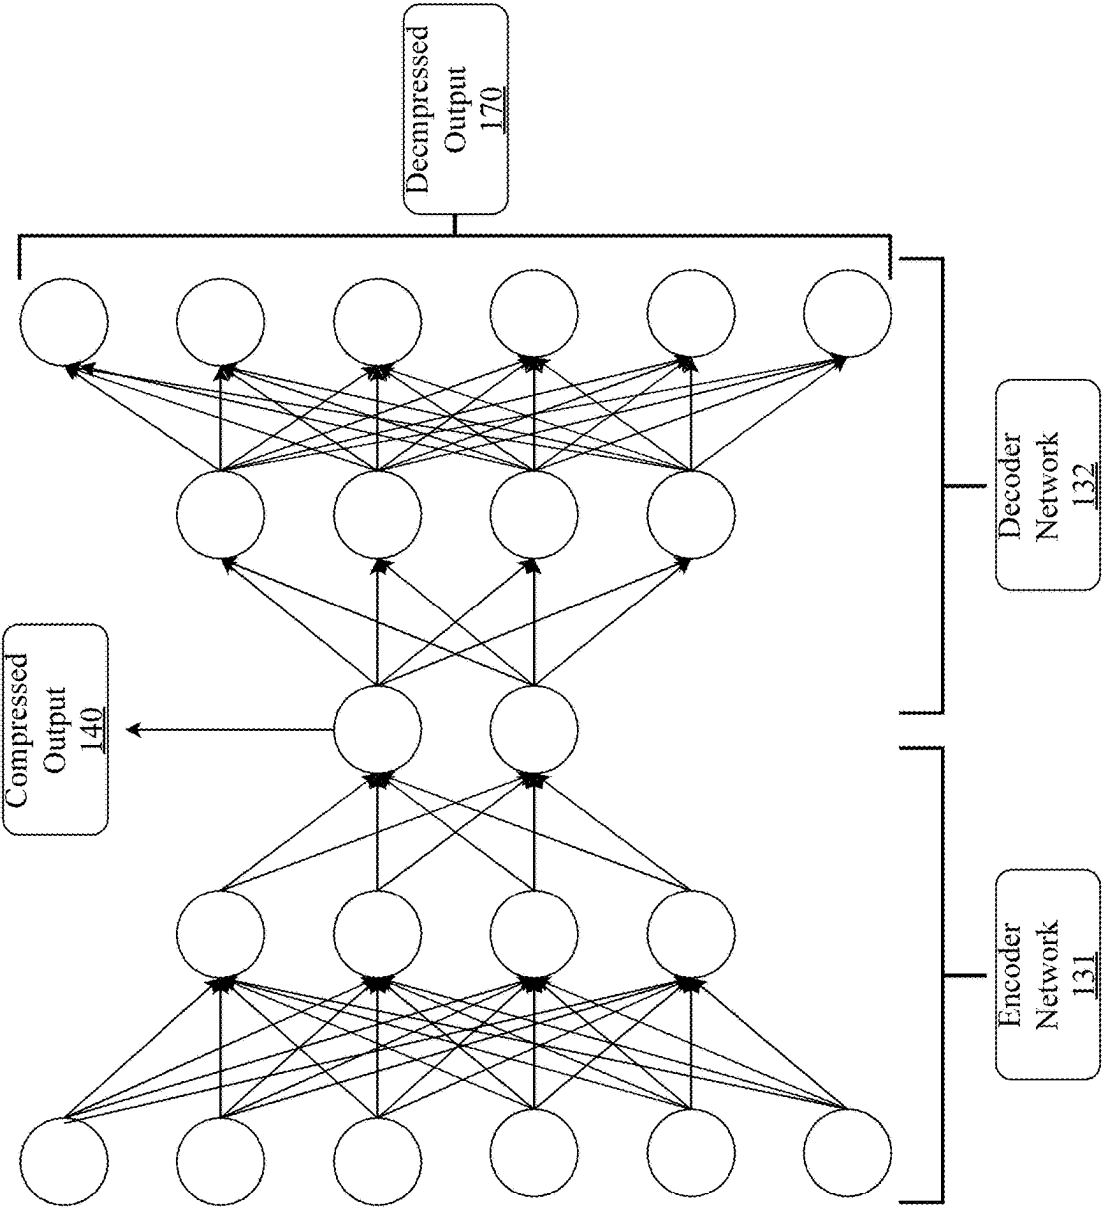


FIG. 2




FIG. 3

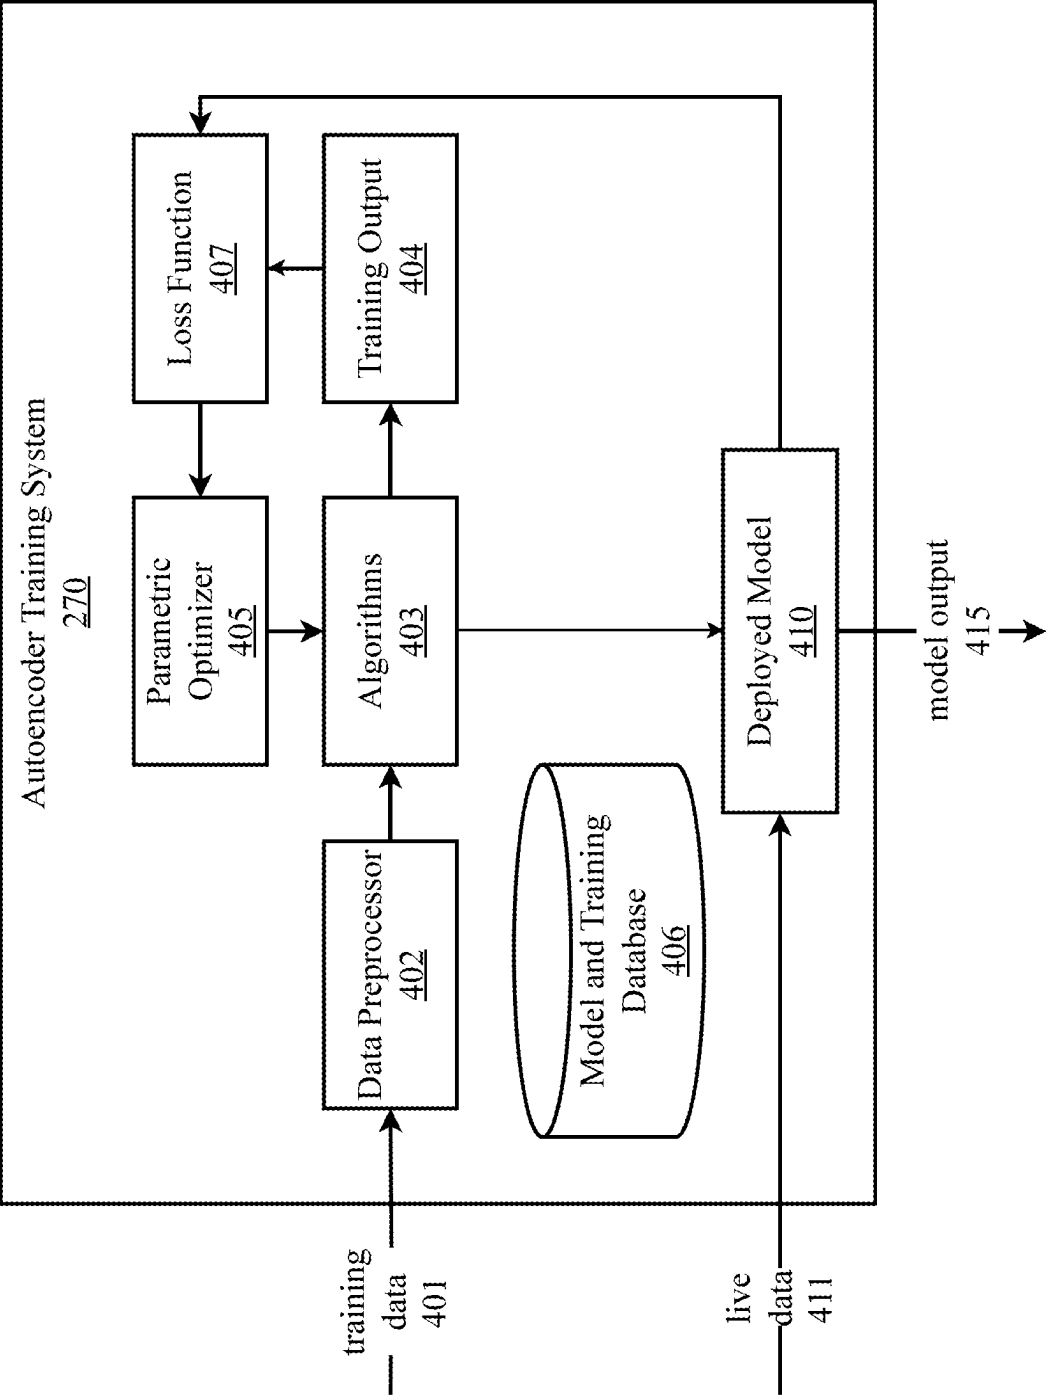


FIG. 4

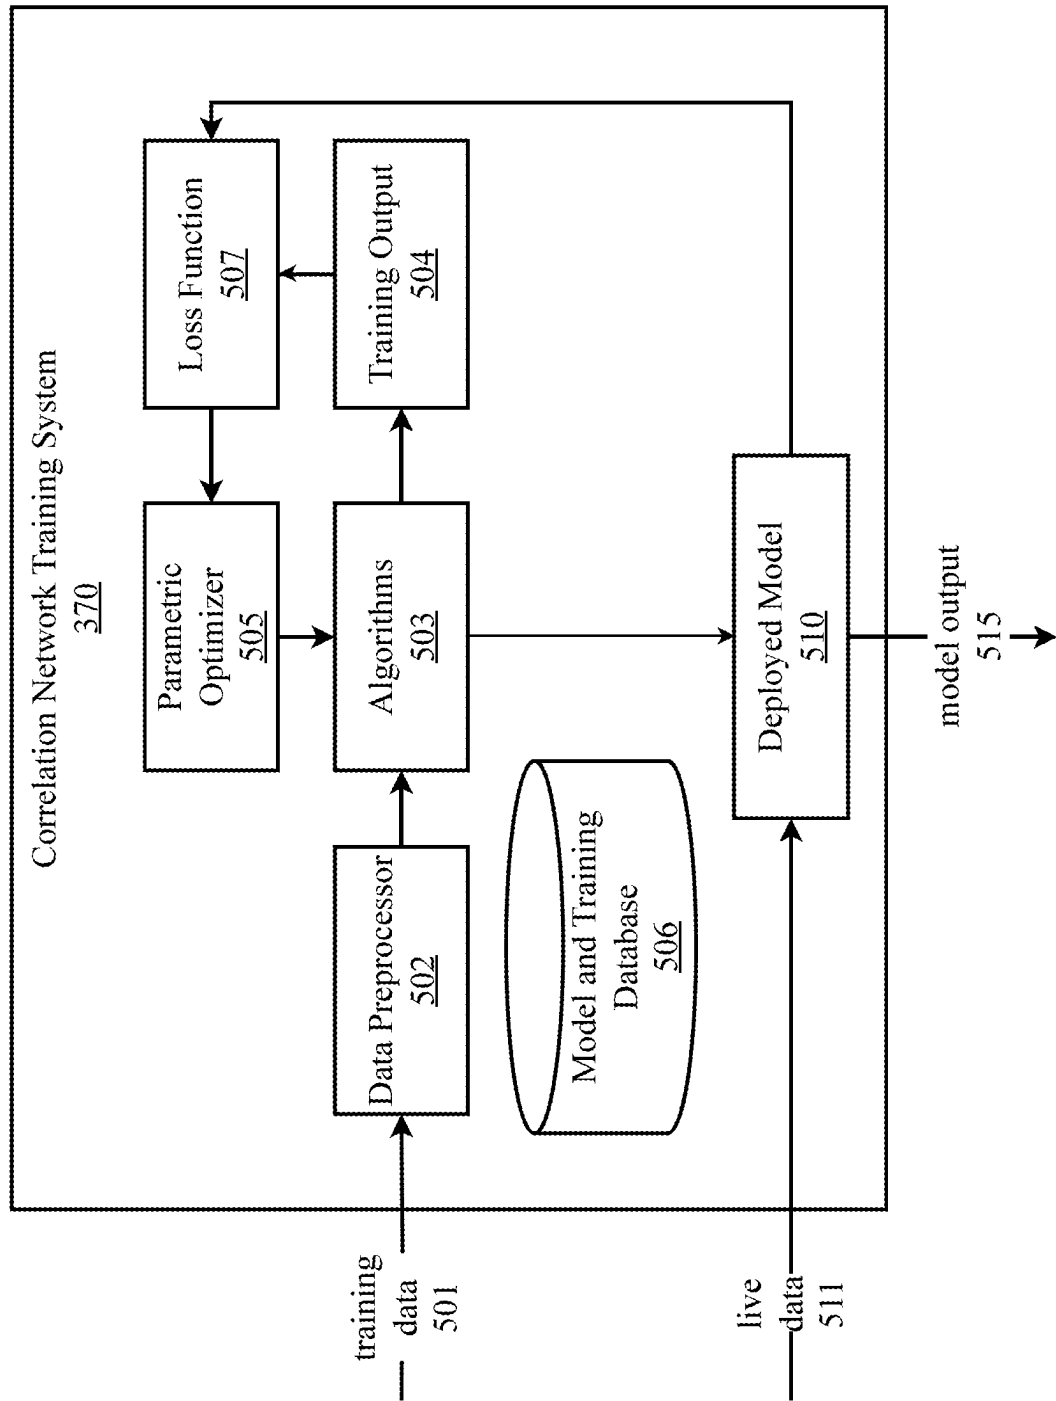


FIG. 5

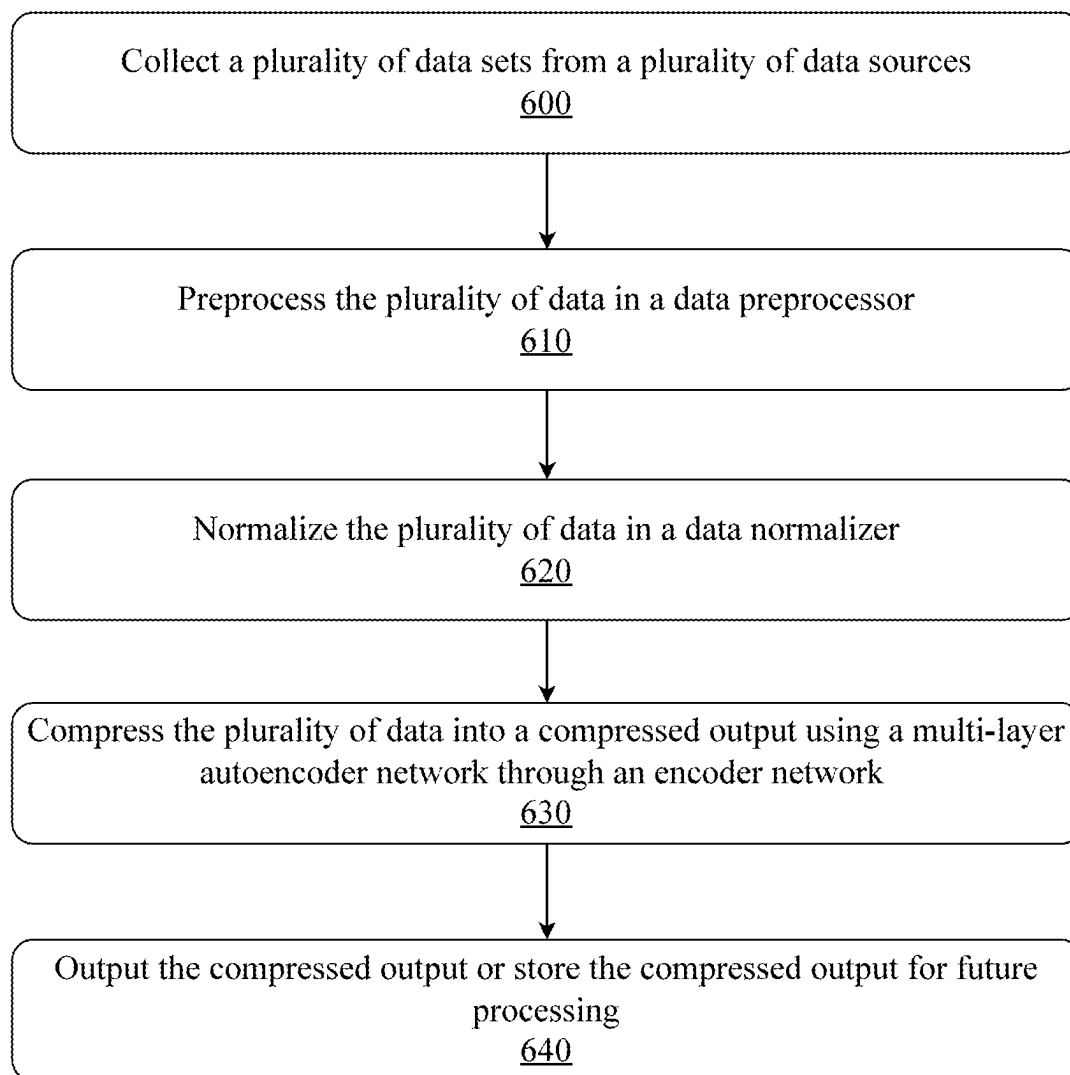


FIG. 6

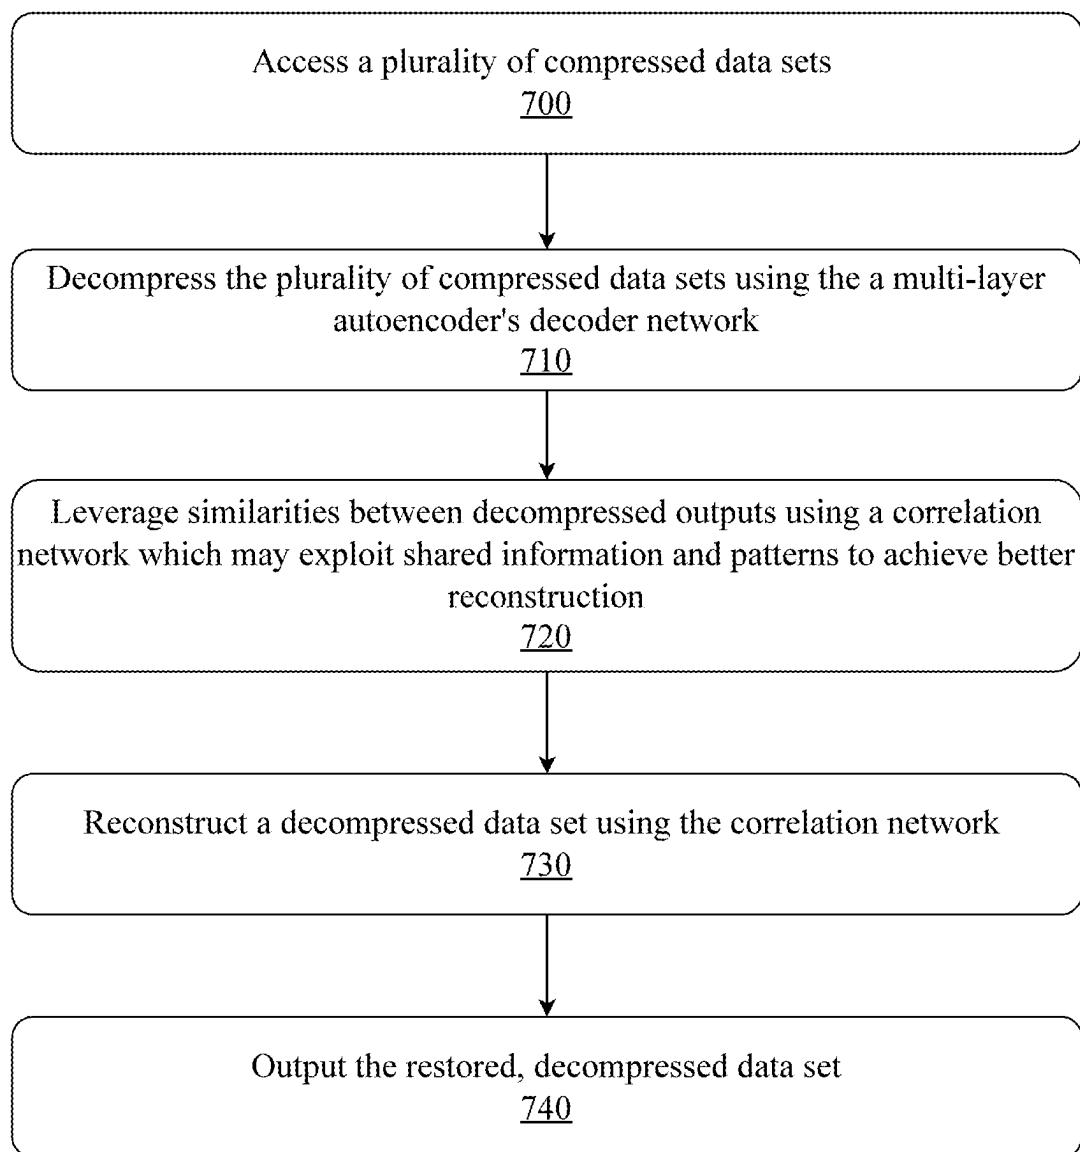


FIG. 7

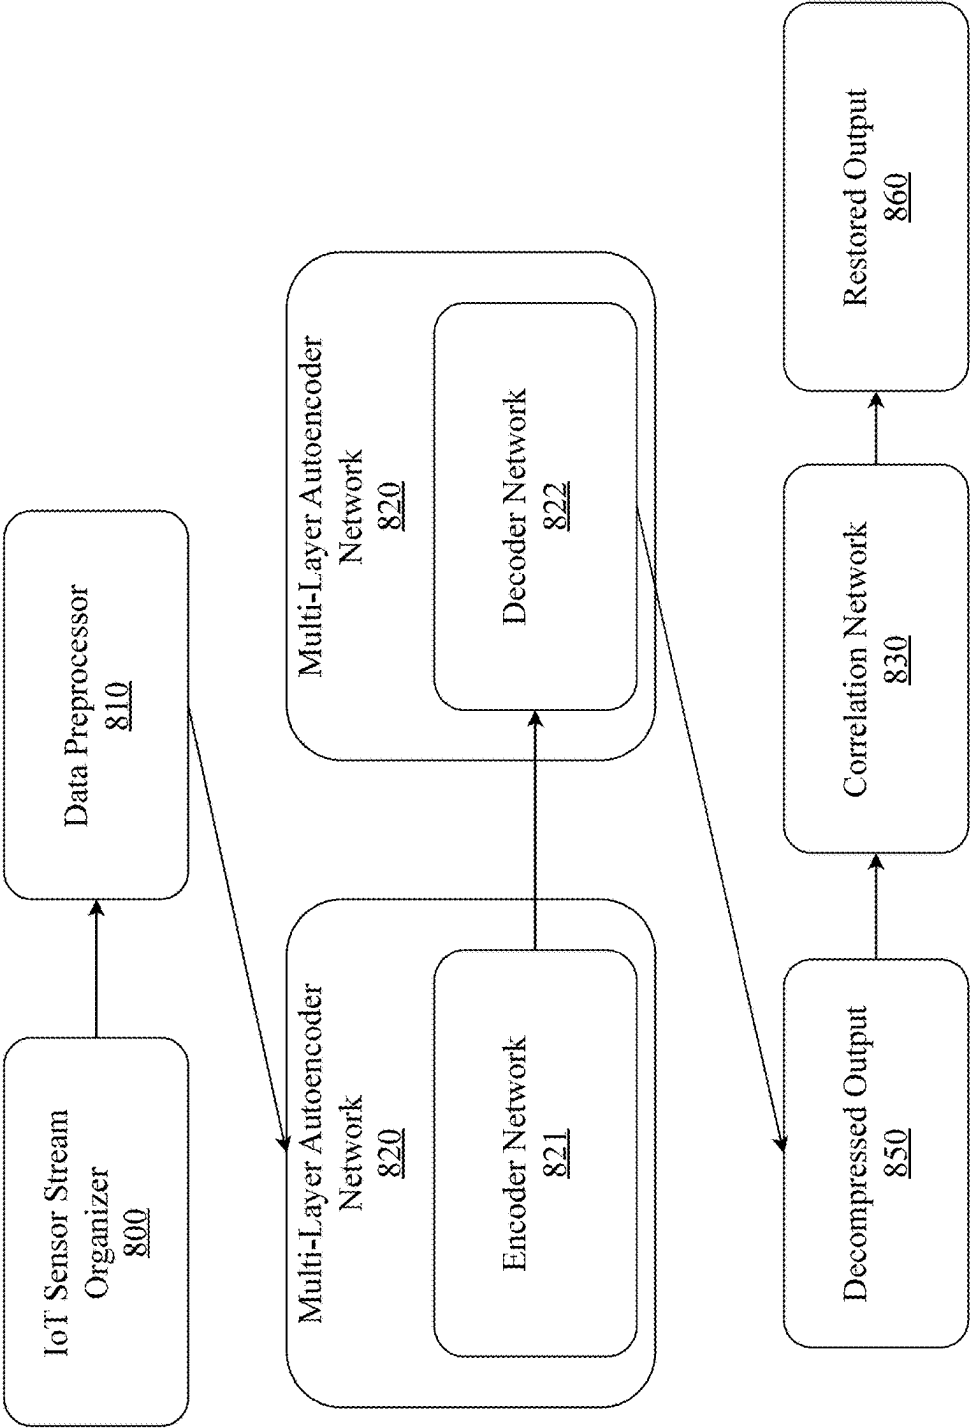


FIG. 8

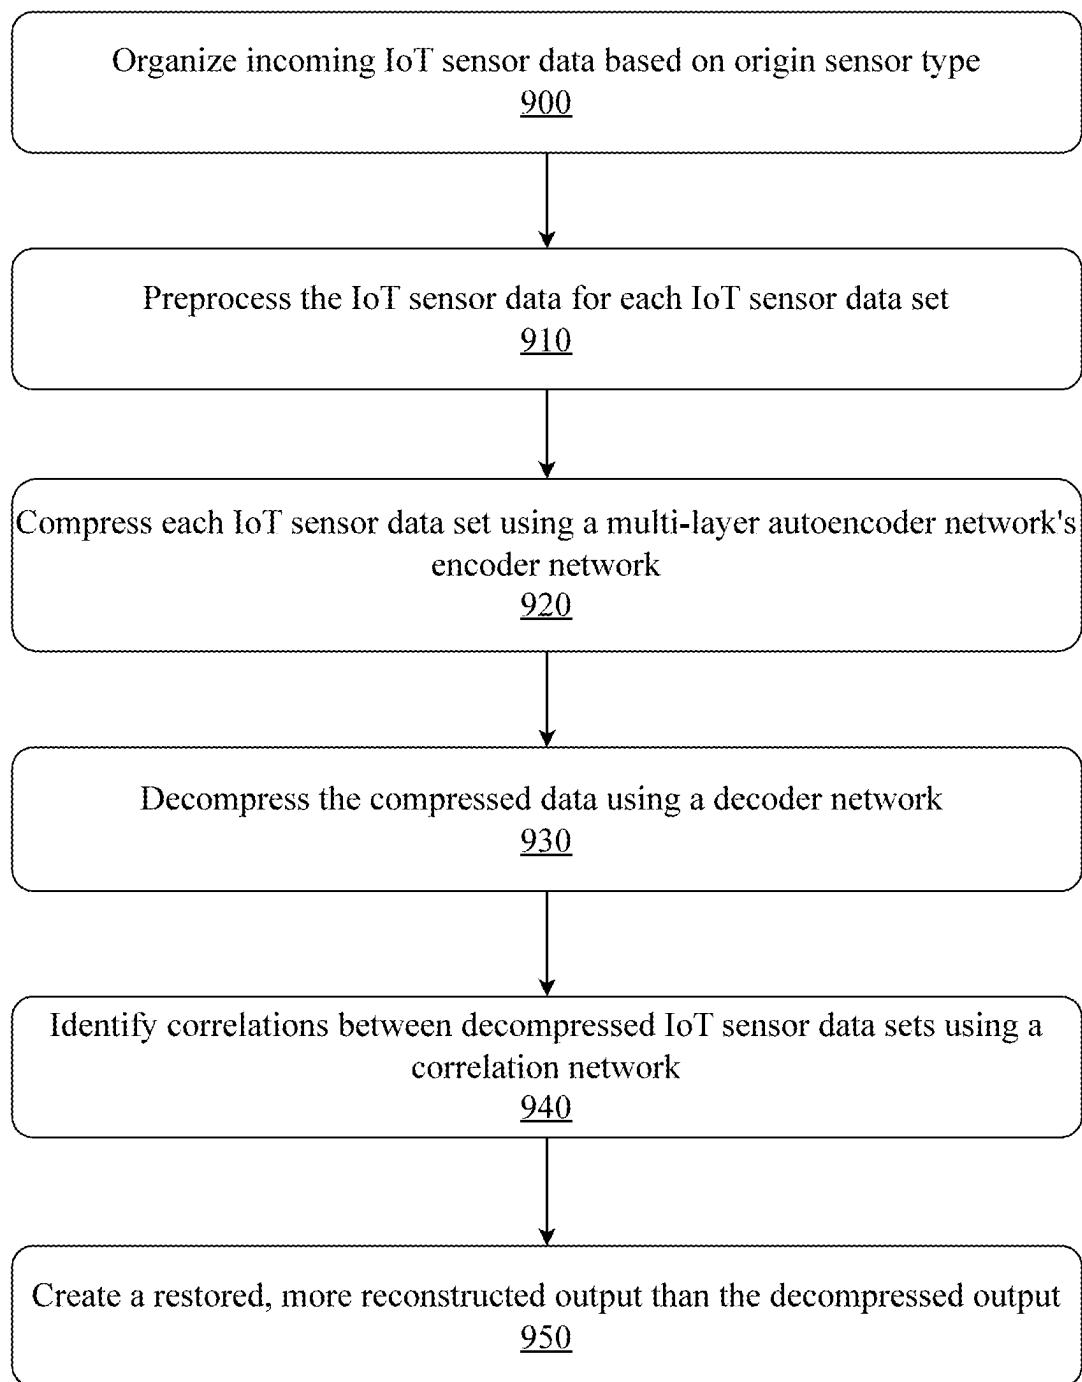


FIG. 9

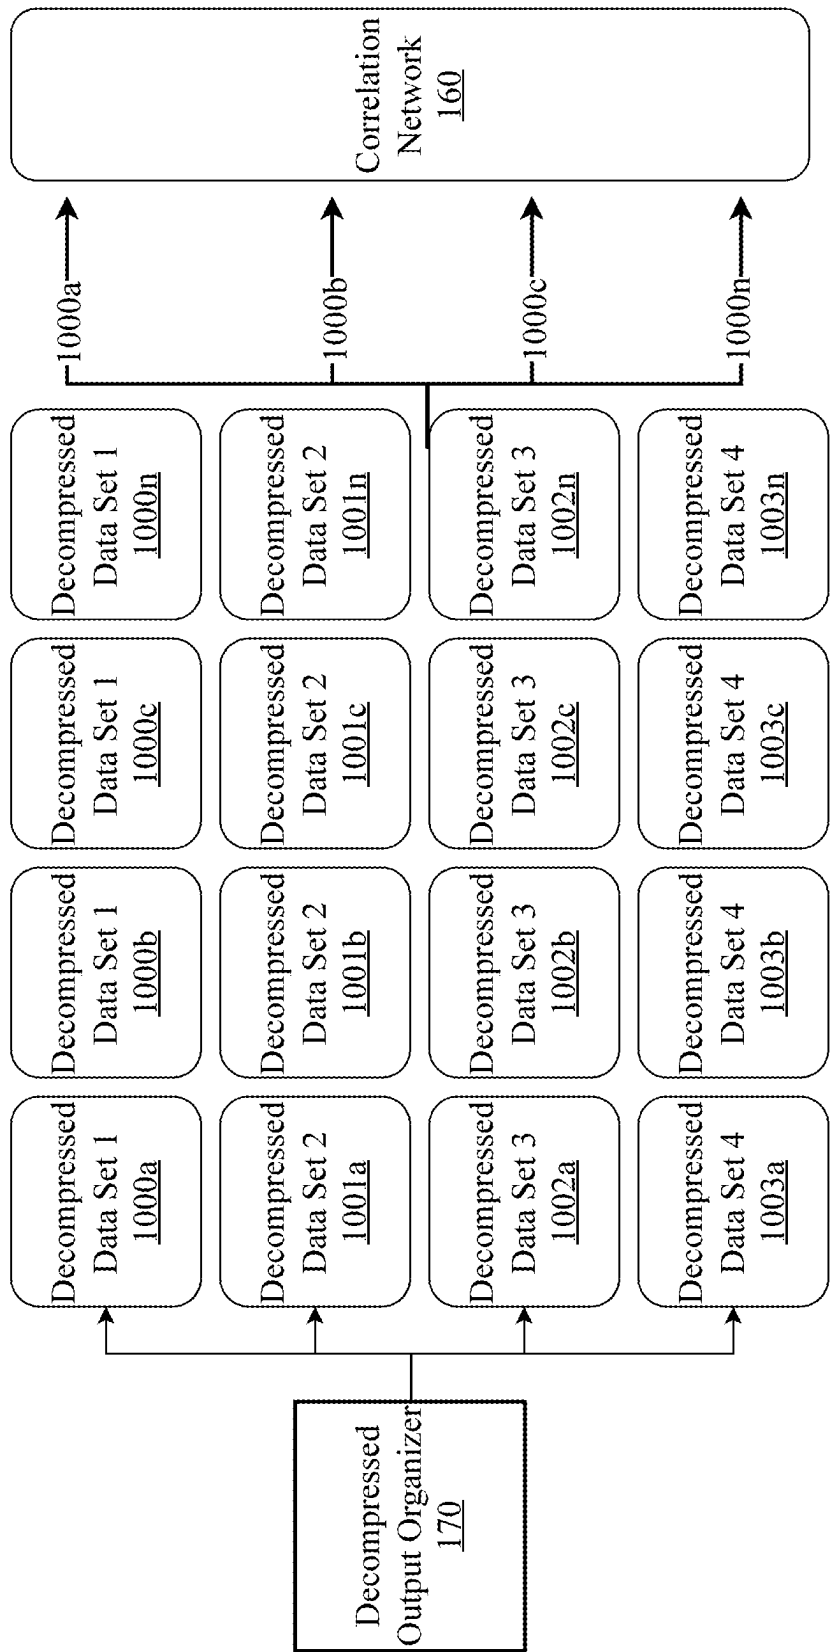


FIG. 10

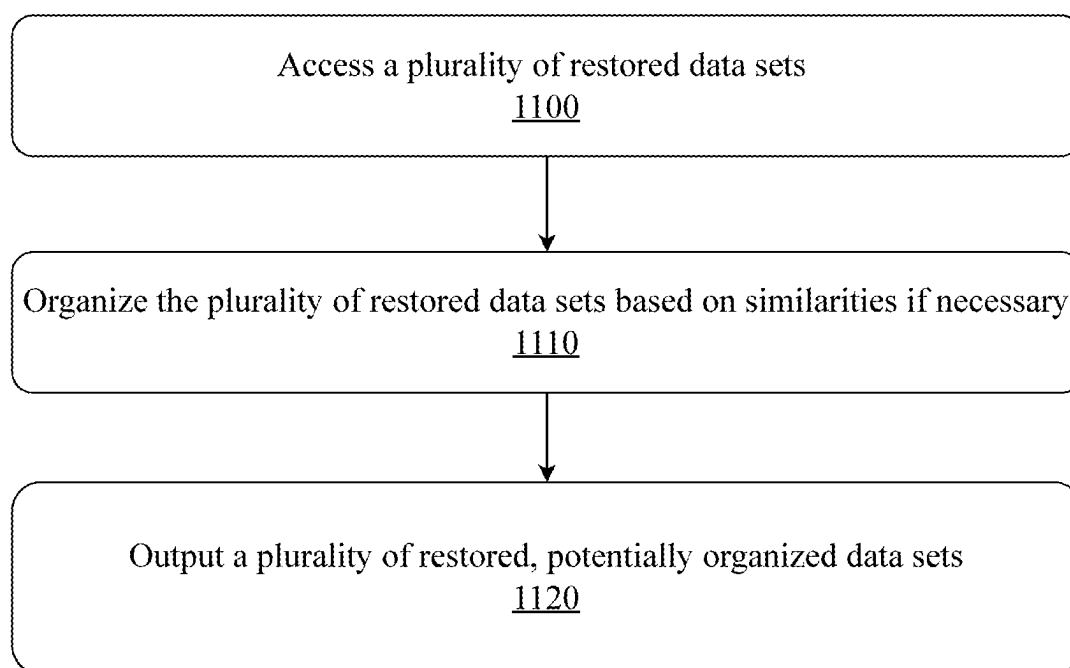


FIG. 11

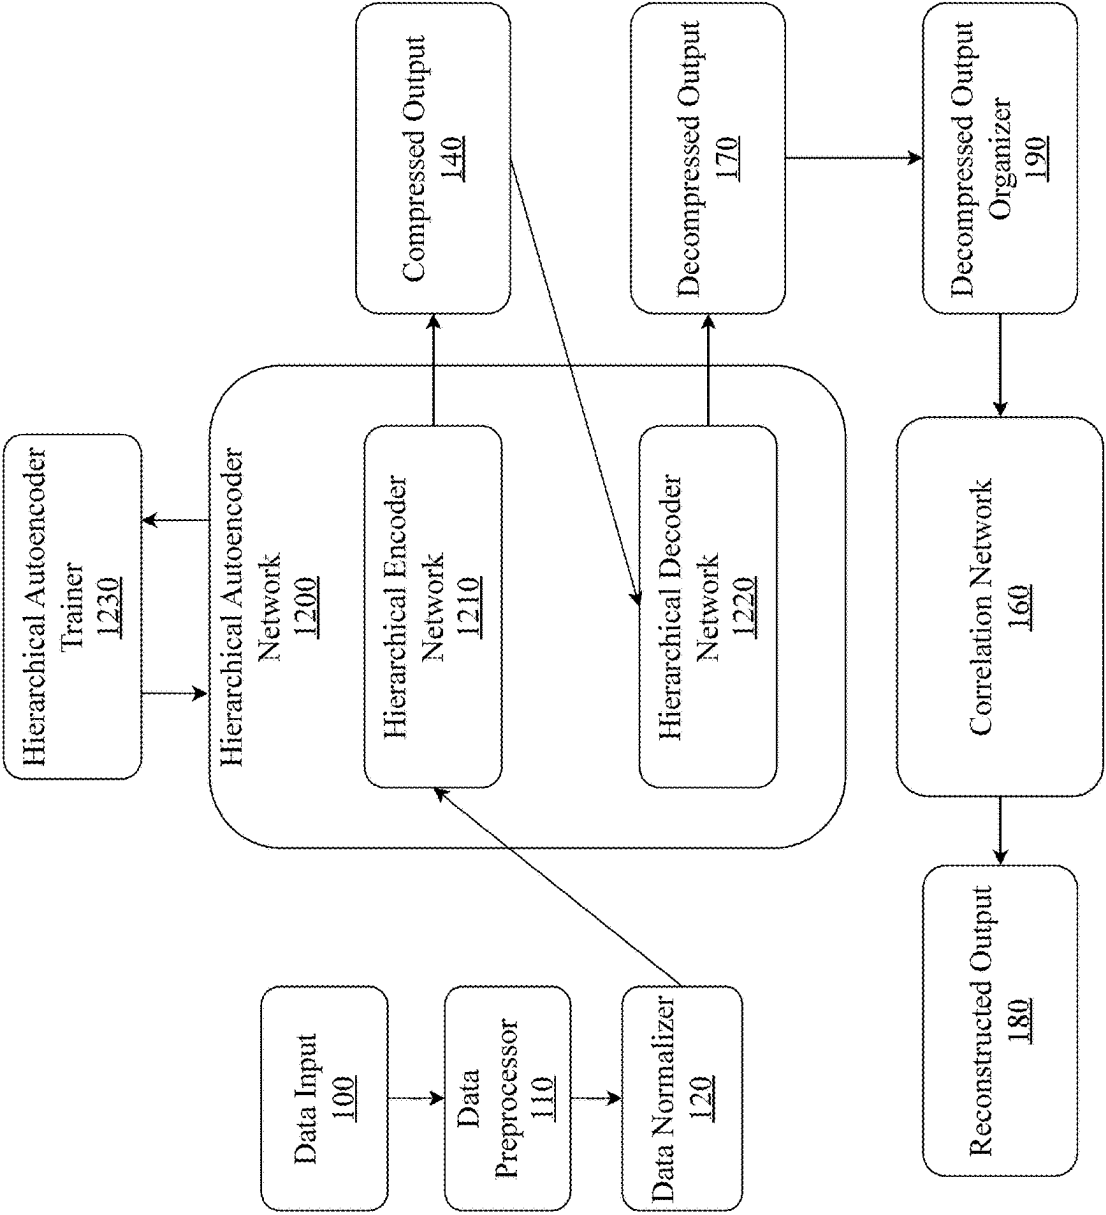


FIG. 12

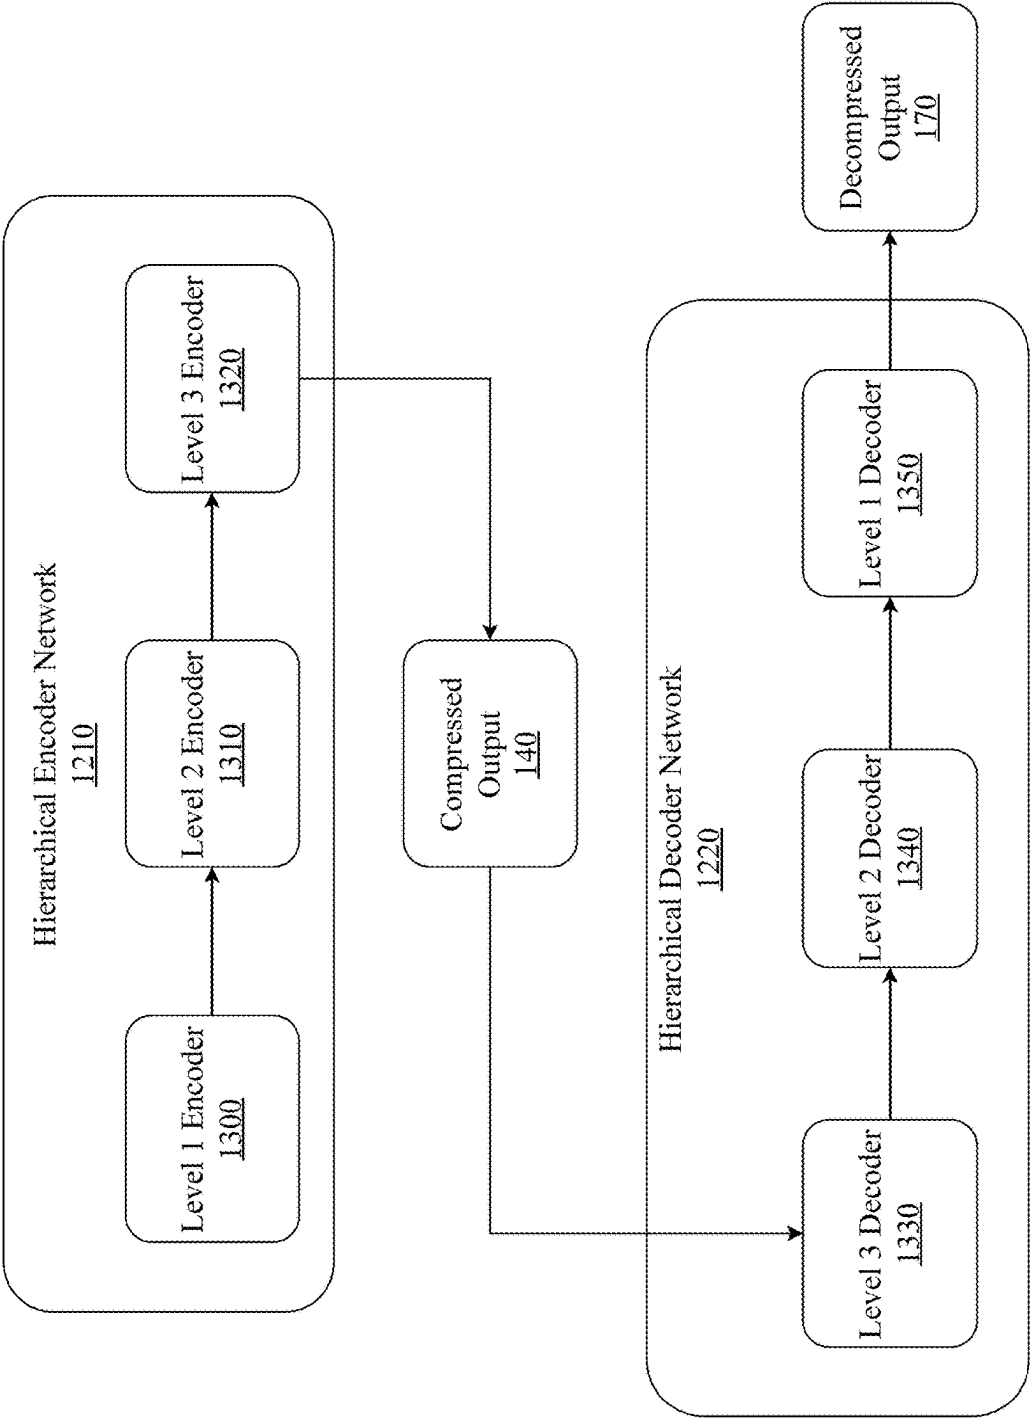


FIG. 13




FIG. 14

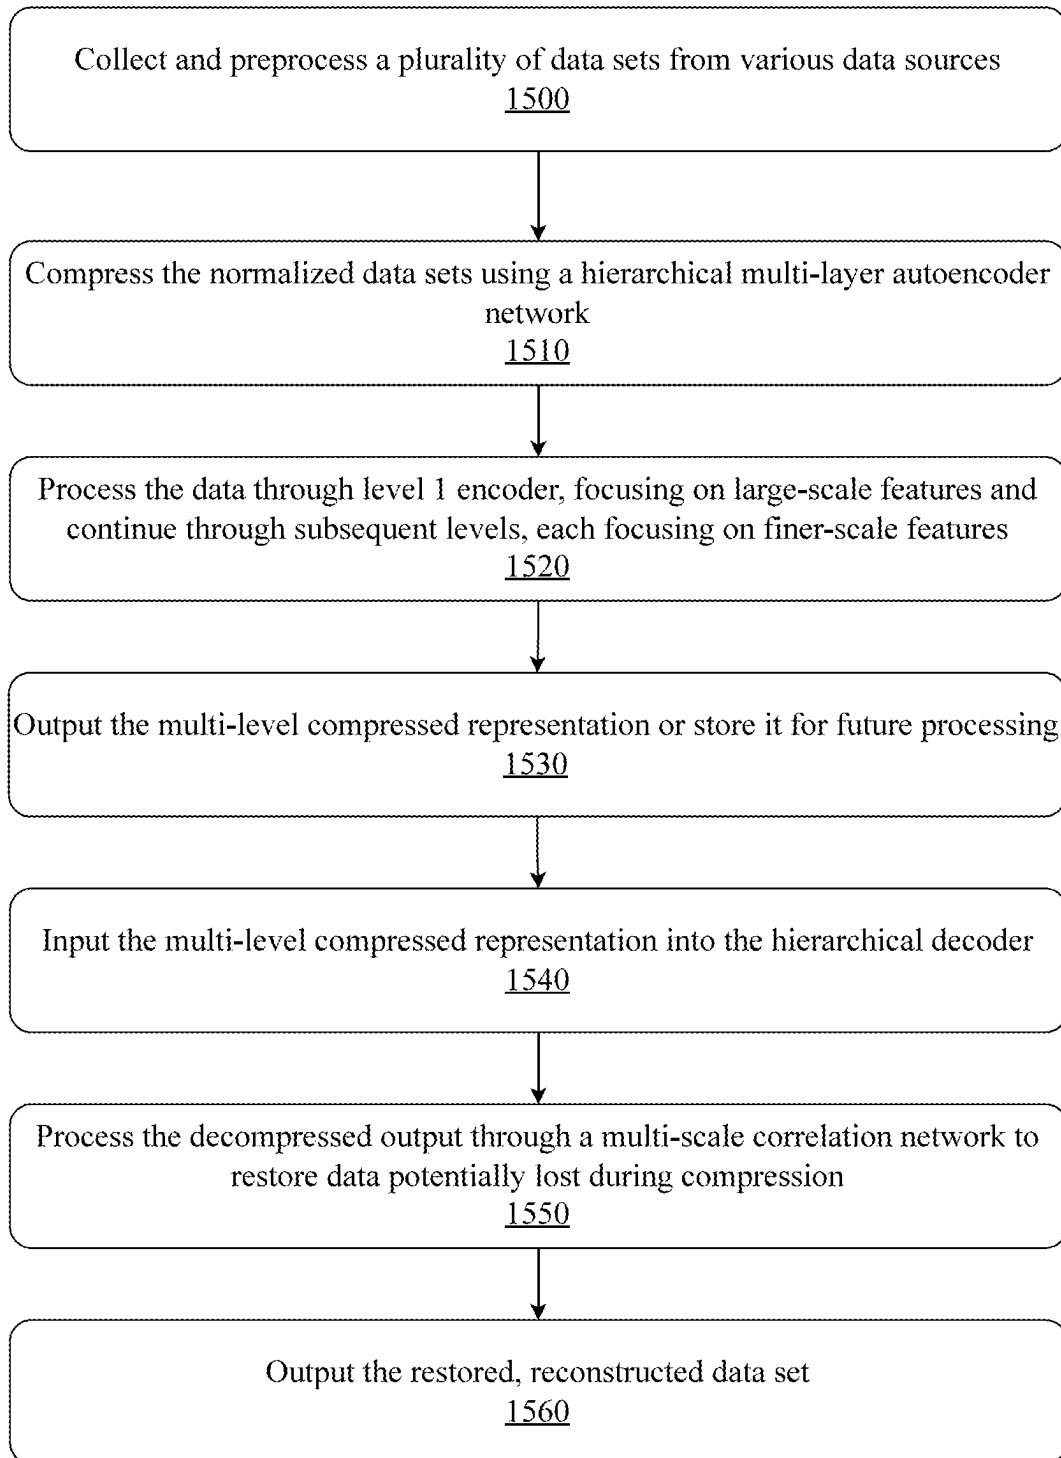


FIG. 15

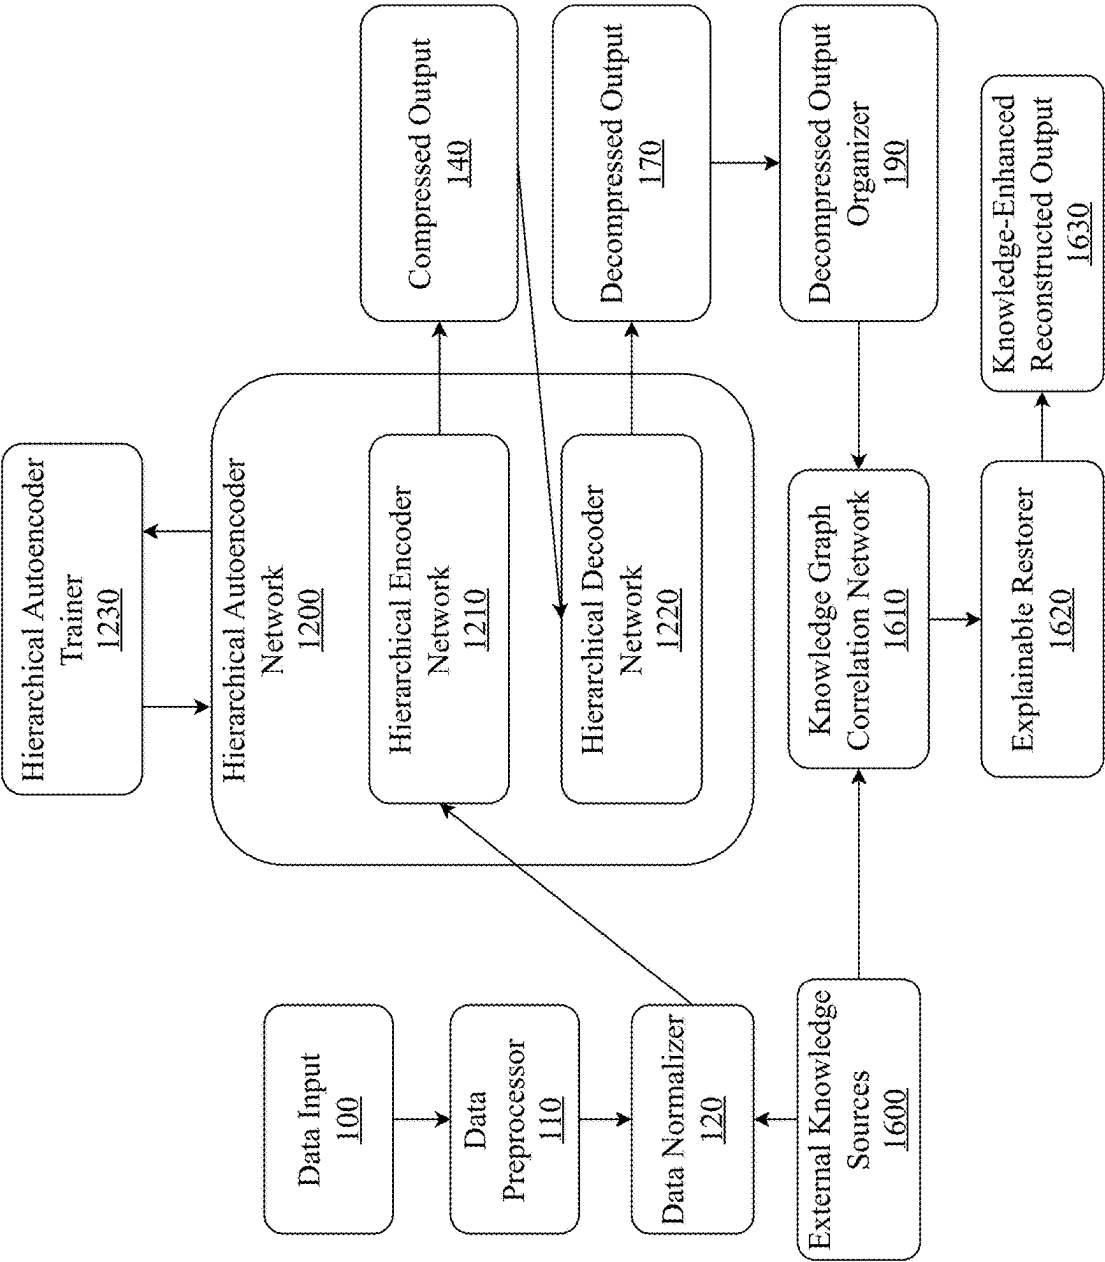


FIG. 16

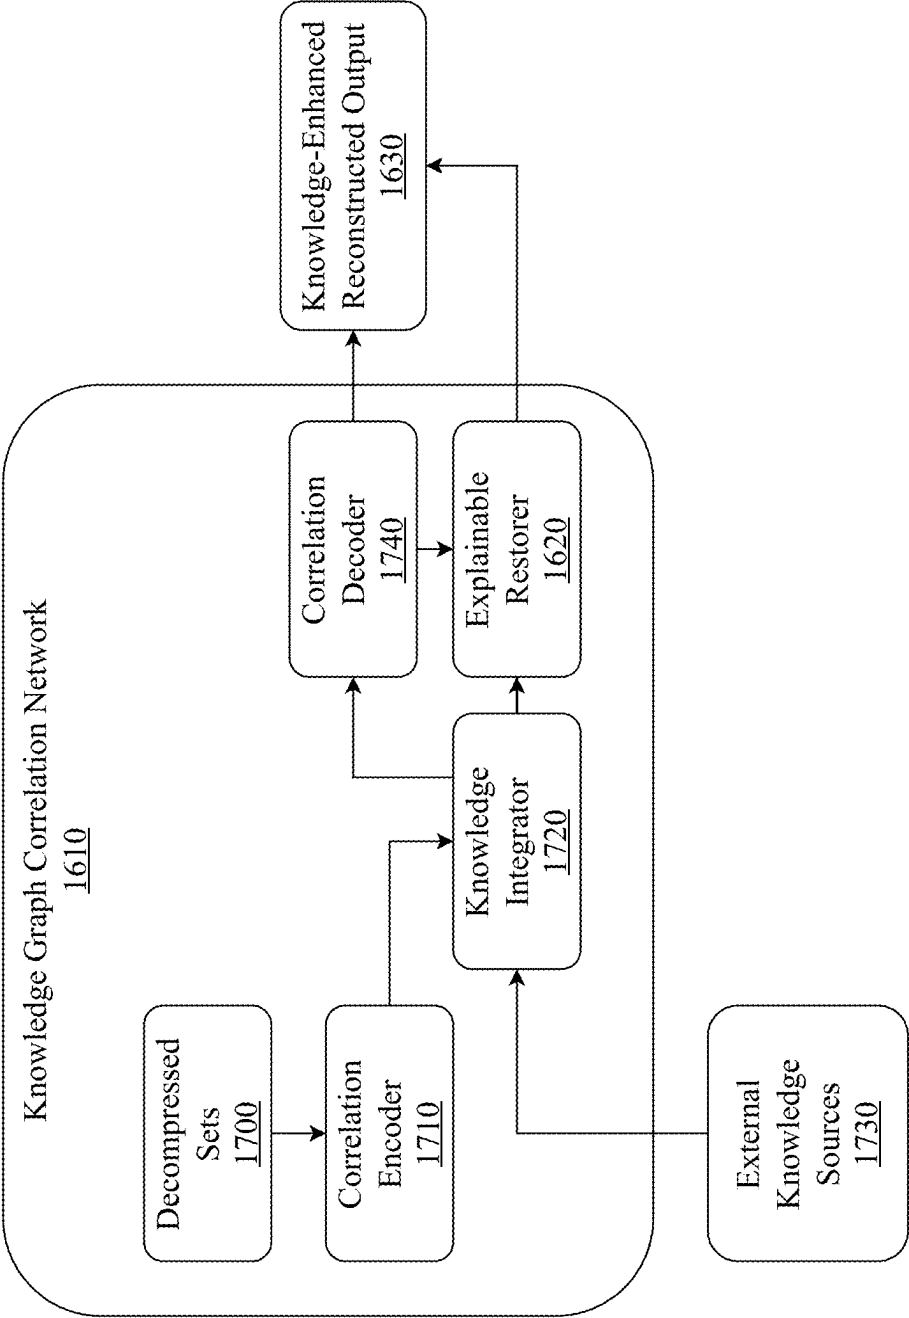


FIG. 17

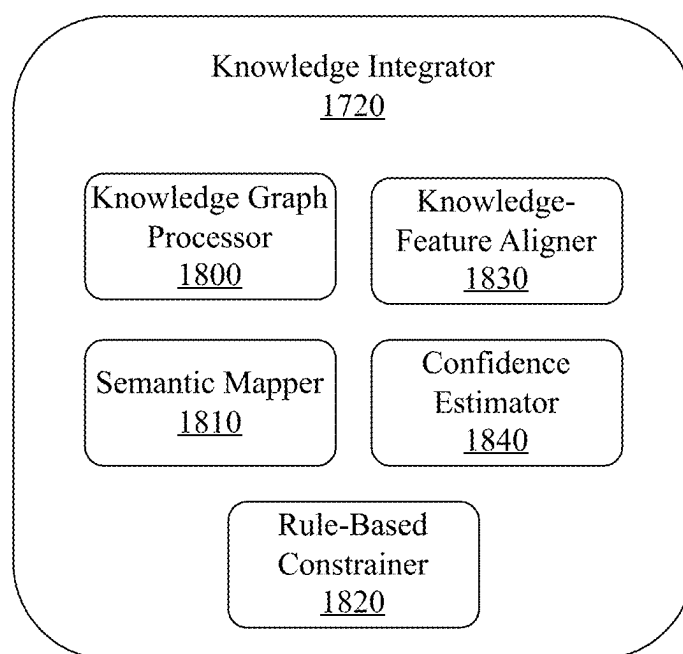


FIG. 18

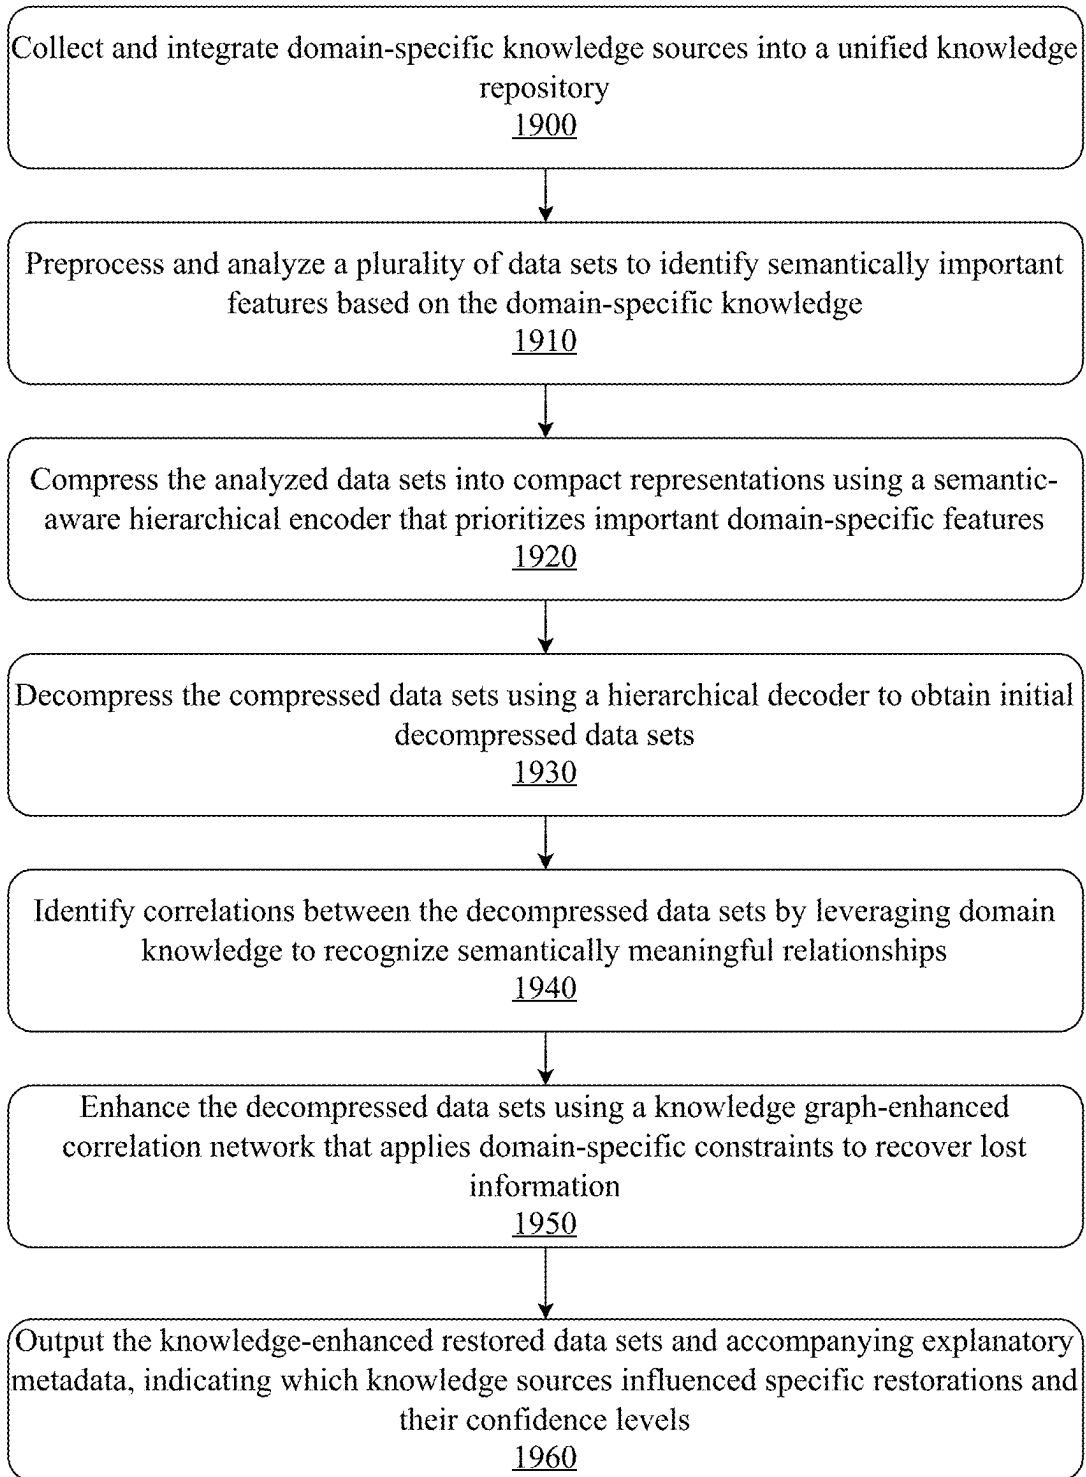


FIG. 19

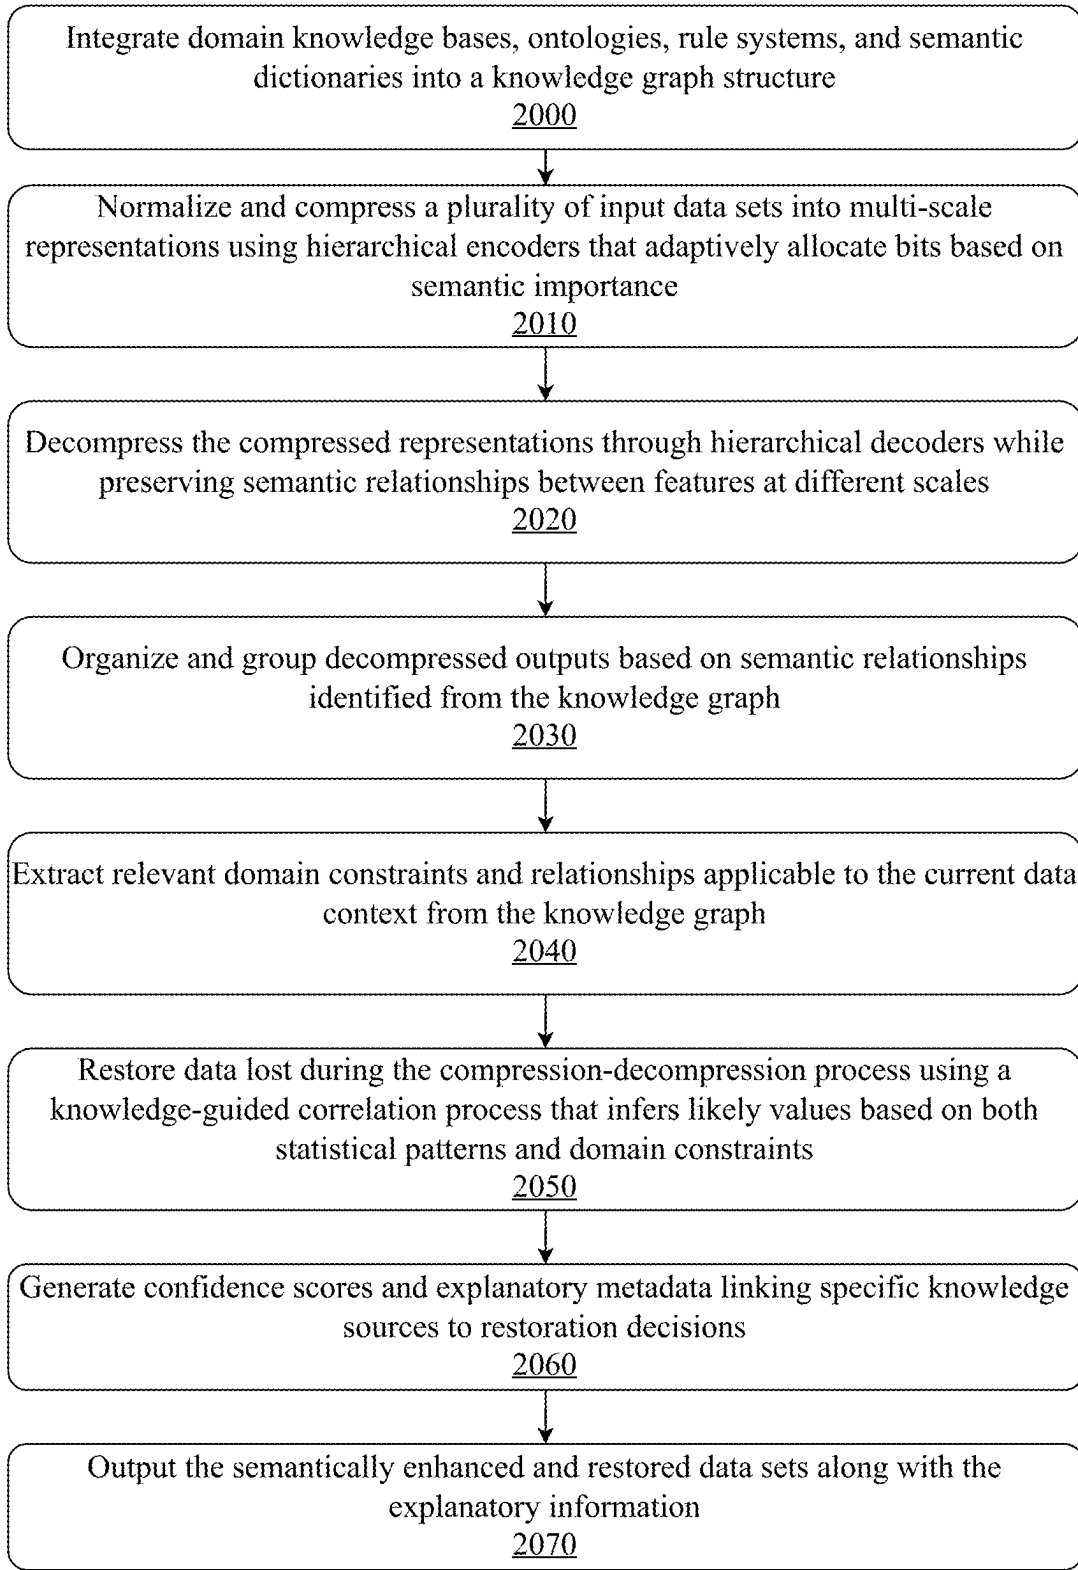
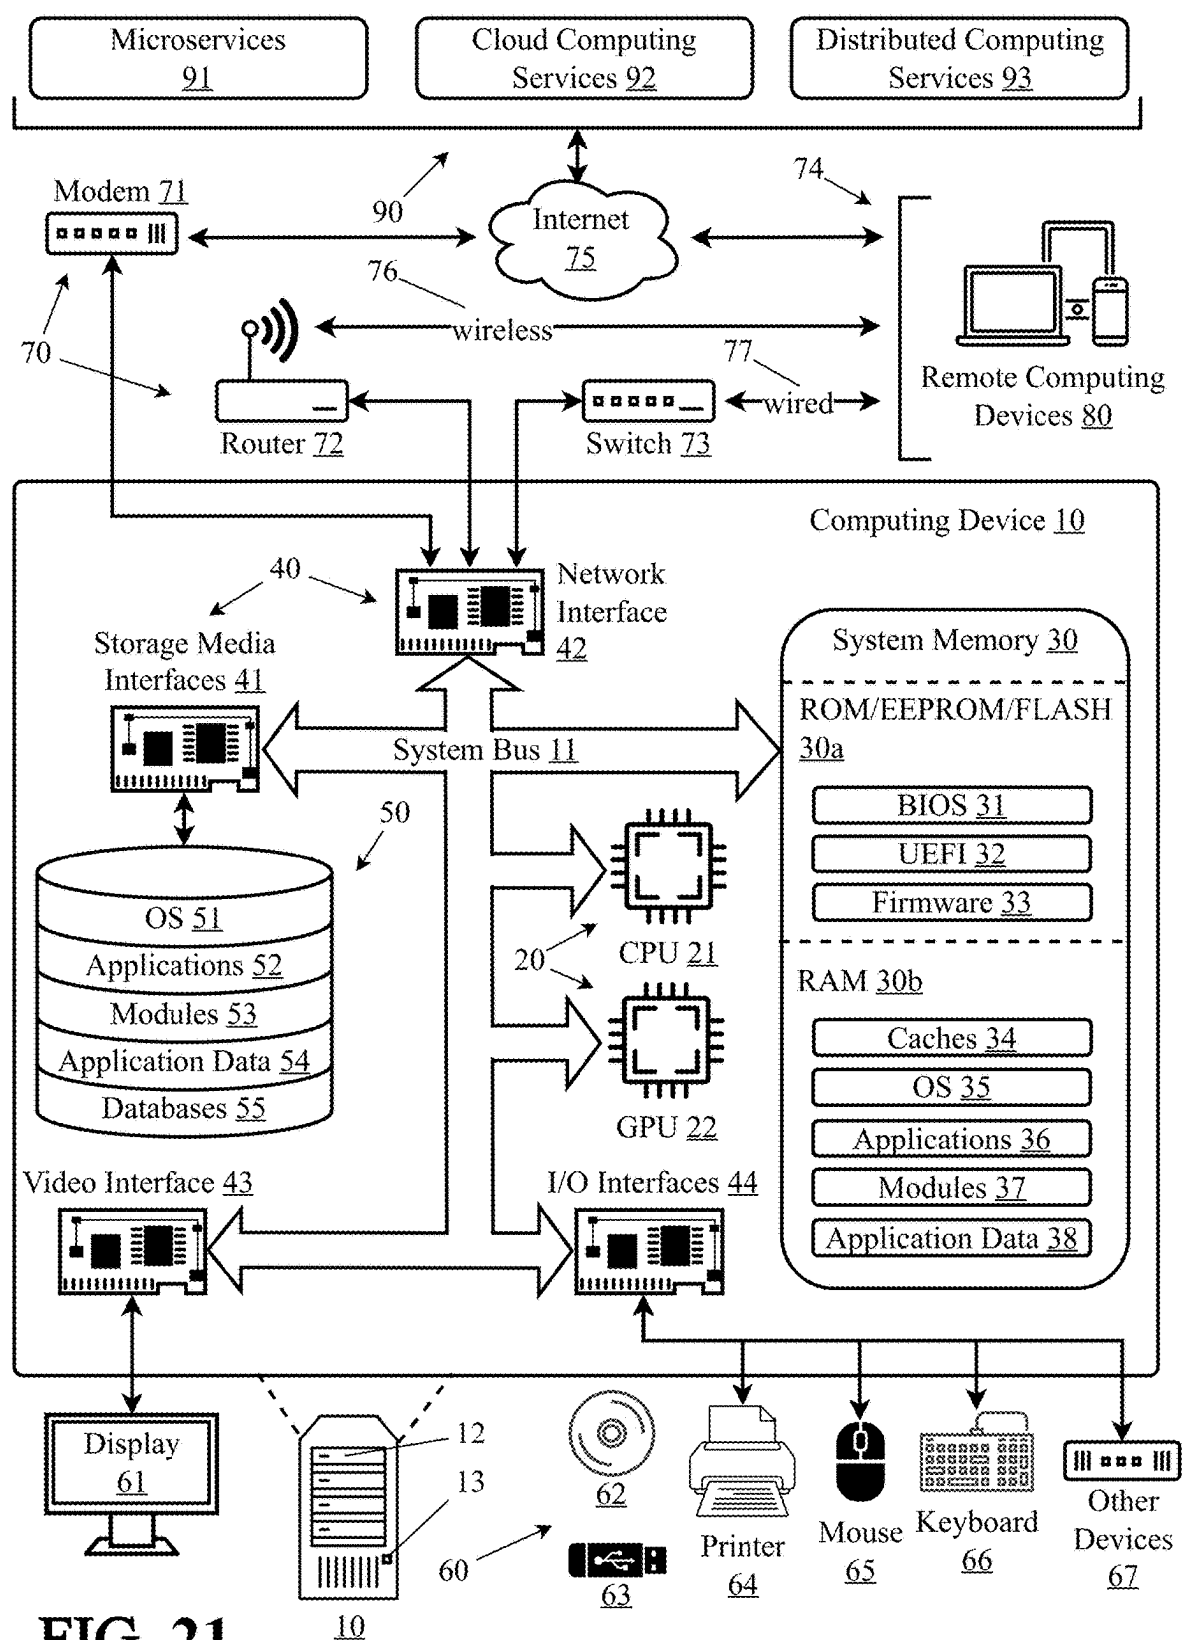


FIG. 20



**SYSTEM AND METHOD FOR
COMPRESSING AND RESTORING DATA
USING HIERARCHICAL AUTOENCODERS
AND A KNOWLEDGE-ENHANCED
CORRELATION NETWORK**

**CROSS-REFERENCE TO RELATED
APPLICATIONS**

[0001] Priority is claimed in the application data sheet to the following patents or patent applications, each of which is expressly incorporated herein by reference in its entirety:

[0002] Ser. No. 18/972,797

[0003] Ser. No. 18/648,340

[0004] Ser. No. 18/427,716

[0005] Ser. No. 18/410,980

[0006] Ser. No. 18/537,728

BACKGROUND OF THE INVENTION

Field of the Art

[0007] The present invention relates to the field of deep learning and data compression. More specifically, the invention pertains to systems and methods that utilize multi-layer autoencoder for data compression and restoration.

Discussion of the State of the Art

[0008] The In recent years, deep learning approaches have shown promising results in data compression and restoration. Autoencoders, a type of neural network architecture, have emerged as a powerful tool for learning compact representations of data. Autoencoders consist of an encoder network that maps input data to a lower-dimensional latent space and a decoder network that reconstructs the original data from the latent representation.

[0009] Multi-layer autoencoders, also known as stacked autoencoders or deep autoencoders, have been proposed to learn hierarchical representations of data. These architectures stack multiple layers of encoders and decoders, allowing for more complex and abstract feature learning. Multi-layer autoencoders have been successfully applied in various domains, such as image compression, video compression, and speech enhancement. However, existing multi-layer autoencoder architectures often focus solely on the compression aspect and do not fully exploit the correlations and patterns within the data. Correlations between different data samples or neighboring regions can provide valuable information for data restoration and enhancement.

[0010] To address this limitation, correlation-based methods have been explored in the context of data restoration. These methods leverage the correlations and similarities between data samples to predict missing or corrupted information. For example, non-local means filtering and block-matching and 3D filtering (BM3D) have been widely used for image denoising by exploiting self-similarities within an image.

[0011] What is needed is an efficient and effective system and method that combine the benefits of multi-layer autoencoders and correlation-based techniques for data compression and restoration. Such an architecture should be capable of learning hierarchical representations, exploiting correlations within the data, and achieving high compression ratios while preserving important information. The present invention addresses these challenges by introducing a multi-layer

autoencoder with a correlation layer. The proposed architecture leverages the power of deep learning to learn compact representations of data while explicitly modeling and utilizing correlations for enhanced data restoration. By incorporating a correlation layer within the autoencoder framework, the invention aims to achieve state-of-the-art performance in data compression and restoration tasks.

SUMMARY OF THE INVENTION

[0012] Accordingly, the inventor has conceived and reduced to practice, a system and method for compressing and restoring data using hierarchical autoencoders and a knowledge-enhanced correlation network. The invention comprises two main components: a multi-level autoencoder for compression and a correlation network for restoration. The multi-level autoencoder consists of an encoder and a decoder. The encoder compresses an input image into a compact representation, while the decoder reconstructs the image from the compressed representation. The architecture of the autoencoder can include convolutional layers, pooling layers, and activation functions. The correlation network may be trained on sets of compressed data sets to learn correlations between them. It takes the compressed representations as input and outputs restored versions of the images. The architecture of the correlation network may include convolutional layers and activation functions. During training, the autoencoder and correlation network may be optimized jointly using loss functions that measure the reconstruction quality and the restoration quality, respectively. The training process may involve forward passes through both networks, calculating losses, and updating the network parameters using optimization algorithms.

[0013] According to a preferred embodiment, a computer system comprising: a hardware memory, wherein the computer system is configured to execute software instructions stored on nontransitory machine-readable storage media that: train a correlation network using sets of cross-correlated training data sets; obtain a plurality of input data sets; compress the input data sets into compressed data sets using a plurality of encoders within a plurality of hierarchical autoencoders; decompress the compressed data sets using a plurality of decoders within the plurality of hierarchical autoencoders to obtain decompressed data sets; and enhance the decompressed data sets using a restoration process that incorporates external information sources to recover lost or degraded information, thereby generating enhanced restored data sets, is disclosed.

[0014] According to another preferred embodiment, a method for compressing and restoring data, comprising the steps of: training a correlation network using sets of cross-correlated training data sets; obtaining a plurality of input data sets; compressing the input data sets into compressed data sets using a plurality of encoders within a plurality of hierarchical autoencoders; decompressing the compressed data sets using a plurality of decoders within the plurality of hierarchical autoencoders to obtain decompressed data sets; and enhancing the decompressed data sets using a restoration process that incorporates external information sources to recover lost or degraded information, thereby generating enhanced restored data sets, is disclosed.

[0015] According to an aspect of an embodiment, the plurality of data sets include hyperspectral data.

[0016] According to an aspect of an embodiment, the data sets comprise data organized by type prior to preprocessing.

[0017] According to an aspect of an embodiment, the restoration process utilizes a correlation network that leverages semantic relationships within the data to enhance restoration quality.

BRIEF DESCRIPTION OF THE DRAWING FIGURES

[0018] FIG. 1 is a block diagram illustrating an exemplary system architecture for compressing and restoring data using multi-level autoencoders and correlation networks.

[0019] FIG. 2 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, an autoencoder network.

[0020] FIG. 3 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a correlation network.

[0021] FIG. 4 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, an autoencoder training system.

[0022] FIG. 5 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, correlation network training system.

[0023] FIG. 6 is a flow diagram illustrating an exemplary method for compressing a data input using a system for compressing and restoring data using multi-level autoencoders and correlation networks.

[0024] FIG. 7 is a flow diagram illustrating an exemplary method for decompressing a compressed data input using system for compressing and restoring data using multi-level autoencoders and correlation networks.

[0025] FIG. 8 is a block diagram illustrating an exemplary system architecture for compressing and restoring IoT sensor data using a system for compressing and restoring data using multi-level autoencoders and correlation networks.

[0026] FIG. 9 is a flow diagram illustrating an exemplary method for compressing and decompressing IoT sensor data using a system for compressing and restoring data using multi-level autoencoders and correlation networks.

[0027] FIG. 10 is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, the decompressed output organizer.

[0028] FIG. 11 is a flow diagram illustrating an exemplary method for organizing restored, decompressed data sets after correlation network processing.

[0029] FIG. 12 is a block diagram illustrating an exemplary system architecture for compressing and restoring data using hierarchical autoencoders and correlation networks.

[0030] FIG. 13 is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using hierarchical autoencoders and correlation networks, a hierarchical autoencoder.

[0031] FIG. 14 is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using hierarchical autoencoders and correlation networks, a hierarchical autoencoder trainer.

[0032] FIG. 15 is a flow diagram illustrating an exemplary method for compressing and restoring data using hierarchical autoencoders and correlation networks.

[0033] FIG. 16 is a block diagram of an exemplary system architecture for compressing and restoring data using hierarchical autoencoders and knowledge-enhanced correlation networks.

[0034] FIG. 17 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data, a knowledge graph correlation network.

[0035] FIG. 18 is a block diagram illustrating an exemplary architecture for a subsystem of the knowledge graph correlation network, a knowledge integrator.

[0036] FIG. 19 is a flow diagram illustrating an exemplary method for knowledge-enhanced data compression and restoration using hierarchical autoencoders and correlation networks.

[0037] FIG. 20 is a flow diagram illustrating an alternative exemplary method for semantically guided multi-level data compression and knowledge-based restoration using hierarchical autoencoders and knowledge graph correlation networks.

[0038] FIG. 21 illustrates an exemplary computing environment on which an embodiment described herein may be implemented, in full or in part.

DETAILED DESCRIPTION OF THE INVENTION

[0039] The inventor has conceived, and reduced to practice, a system and method for compressing and restoring data using hierarchical autoencoders and a knowledge-enhanced correlation network.

[0040] One or more different aspects may be described in the present application. Further, for one or more of the aspects described herein, numerous alternative arrangements may be described; it should be appreciated that these are presented for illustrative purposes only and are not limiting of the aspects contained herein or the claims presented herein in any way. One or more of the arrangements may be widely applicable to numerous aspects, as may be readily apparent from the disclosure. In general, arrangements are described in sufficient detail to enable those skilled in the art to practice one or more of the aspects, and it should be appreciated that other arrangements may be utilized and that structural, logical, software, electrical and other changes may be made without departing from the scope of the particular aspects. Particular features of one or more of the aspects described herein may be described with reference to one or more particular aspects or figures that form a part of the present disclosure, and in which are shown, by way of illustration, specific arrangements of one or more of the aspects. It should be appreciated, however, that such features are not limited to usage in the one or more particular aspects or figures with reference to which they are described. The present disclosure is neither a literal description of all arrangements of one or more of the aspects nor a listing of features of one or more of the aspects that must be present in all arrangements.

[0041] Headings of sections provided in this patent application and the title of this patent application are for convenience only, and are not to be taken as limiting the disclosure in any way.

[0042] Devices that are in communication with each other need not be in continuous communication with each other, unless expressly specified otherwise. In addition, devices that are in communication with each other may communicate directly or indirectly through one or more communication means or intermediaries, logical or physical.

[0043] A description of an aspect with several components in communication with each other does not imply that all such components are required. To the contrary, a variety of optional components may be described to illustrate a wide variety of possible aspects and in order to more fully illustrate one or more aspects. Similarly, although process steps, method steps, algorithms or the like may be described in a sequential order, such processes, methods and algorithms may generally be configured to work in alternate orders, unless specifically stated to the contrary. In other words, any sequence or order of steps that may be described in this patent application does not, in and of itself, indicate a requirement that the steps be performed in that order. The steps of described processes may be performed in any order practical. Further, some steps may be performed simultaneously despite being described or implied as occurring non-simultaneously (e.g., because one step is described after the other step). Moreover, the illustration of a process by its depiction in a drawing does not imply that the illustrated process is exclusive of other variations and modifications thereto, does not imply that the illustrated process or any of its steps are necessary to one or more of the aspects, and does not imply that the illustrated process is preferred. Also, steps are generally described once per aspect, but this does not mean they must occur once, or that they may only occur once each time a process, method, or algorithm is carried out or executed. Some steps may be omitted in some aspects or some occurrences, or some steps may be executed more than once in a given aspect or occurrence.

[0044] When a single device or article is described herein, it will be readily apparent that more than one device or article may be used in place of a single device or article. Similarly, where more than one device or article is described herein, it will be readily apparent that a single device or article may be used in place of the more than one device or article. The functionality or the features of a device may be alternatively embodied by one or more other devices that are not explicitly described as having such functionality or features. Thus, other aspects need not include the device itself.

[0045] Techniques and mechanisms described or referenced herein will sometimes be described in singular form for clarity. However, it should be appreciated that particular aspects may include multiple iterations of a technique or multiple instantiations of a mechanism unless noted otherwise. Process descriptions or blocks in figures should be understood as representing modules, segments, or portions of code which include one or more executable instructions for implementing specific logical functions or steps in the process. Alternate implementations are included within the scope of various aspects in which, for example, functions may be executed out of order from that shown or discussed, including substantially concurrently or in reverse order, depending on the functionality involved, as would be understood by those having ordinary skill in the art.

Conceptual Architecture

[0046] FIG. 1 is a block diagram illustrating an exemplary system architecture for compressing and restoring data using multi-level autoencoders and correlation networks. In one embodiment, a system for compressing and restoring data using multi-level autoencoders and correlation networks comprises a plurality of data inputs **100**, a data preprocessor **110**, a data normalizer **120**, a multi-layer autoencoder network **130** which further comprises an encoder network **131** and a decoder network **132**, a plurality of compressed outputs **140**, plurality of decompressed outputs **170**, a decompressed output organizer **190**, a plurality of correlation networks **160**, and a reconstructed output **180**. The plurality of data inputs **100** are representations of raw data from various sources, such as sensors, cameras, or databases. The raw data can be in different formats, including but not limited to images, videos, audio, or structured data. The plurality of data inputs **100** may be transferred to the data preprocessor **110** for further processing. The data preprocessor **110** applies various preprocessing techniques to the raw data received from the data input **100**. These techniques may include data cleaning, noise reduction, artifact removal, or format conversion. The preprocessor **110** ensures that the data is in a suitable format and quality for subsequent stages of the system.

[0047] The preprocessed data may then be passed to the data normalizer **120**. The data normalizer **120** scales and normalizes the data to a consistent range, typically between 0 and 1. Normalization helps to improve the training stability and convergence of the autoencoder network **130**. The normalized data is fed into the autoencoder network **130**, which includes both the encoder network **131** and the decoder network **132**. The encoder network **131** is responsible for encoding the input data into a lower-dimensional latent space representation. It consists of multiple layers of encoders that progressively reduce the dimensionality of the data while capturing the most important features and patterns.

[0048] The compressed latent representation obtained from the encoder network **131** is the compressed output **140**. The compressed output **140** has a significantly reduced size compared to the original input data, enabling efficient storage and transmission. The compressed output **140** may be stored in a storage system. A storage system may be any suitable storage medium, such as a database, file system, or cloud storage. Storage systems allow for the efficient management and retrieval of the compressed data as needed. When the compressed data needs to be restored or reconstructed, it may be retrieved from the storage system and passed to the decoder network **132**. Additionally, the compressed data may be directly passed to either the decompression network **132**. The decoder network **132** is responsible for decoding the compressed latent representation back into the original data space by outputting a decompressed output **170**. It consists of multiple layers of decoders that progressively increase the dimensionality of the data, reconstructing the original input.

[0049] The decompressed output **170** from the encoder network **132** may have some loss of information compared to the original input data due to the compression process. To further enhance the quality of the decompressed output, the system may incorporate a correlation network **160**. The correlation network **160** leverages the correlations and patterns between different compressed inputs to restore the

decompressed output more accurately. It learns to capture the relationships and dependencies within the data, allowing for better reconstruction and restoration of the original information. The correlation network **160** takes the decompressed outputs **170** as inputs. It analyzes the correlations and similarities between the data samples and uses this information to refine and enhance the decompressed output. The refined decompressed output from the correlation network **160** is a reconstructed output **180** of the system. The reconstructed output **180** closely resembles the original input data, with minimal loss of information and improved quality compared to the output from the decoder network **132** alone.

[0050] In one embodiment, the correlation network **160** may receive inputs from a decompressed output organizer **190** which that operates on the decompressed outputs **170** obtained from the decoder network **132**. The decompressed output organizer **190** may organize the decompressed outputs **170** into groups based on their correlations and similarities.

[0051] By grouping decompressed outputs **170** based on similarities, the correlation network **160** will more easily be able to identify correlations between decompressed outputs **170**. The correlation network **160** finds patterns and similarities between decompressed outputs **170** to develop a more holistic reconstructed original input. By priming the correlation network **160** with already grouped, similar compressed outputs **170**, the correlation network **160** will be able to generate even more reliable reconstructions. The multi-layer autoencoder network **130** and the correlation network **160** are trained using a large dataset of diverse samples. The training process involves minimizing the reconstruction loss between the original input data and the decompressed output **170**. The system learns to compress the data efficiently while preserving the essential features and patterns.

[0052] FIG. 2 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a multi-layer autoencoder network. The multi-layer autoencoder network comprises an encoder network **131** or a decoder network **132** that work together to encode and decode data effectively. The encoder network **131** and decoder network **132** within the multi-layer autoencoder network **210** is comprised of a plurality of layers that contribute to the encoding and decoding process. These layers include, but are not limited to, convolutional layers, pooling layers, and a bottleneck layer. Some embodiments also include functions that operate on information including but not limited to rectified linear unit functions, sigmoid functions, and skip connections.

[0053] The convolutional layers are responsible for extracting meaningful features from the input data. They apply convolutional operations using learnable filters to capture spatial patterns and hierarchical representations of the data. The convolutional layers can have different numbers of filters, kernel sizes, and strides to capture features at various scales and resolutions. Skip connections are employed to facilitate the flow of information across different layers of the autoencoder. Skip connections allow the output of a layer to be directly added to the output of a subsequent layer, enabling the network to learn residual mappings and mitigate the vanishing gradient problem. Skip connections help in preserving fine-grained details and improving the training stability of the autoencoder.

[0054] Pooling layers are used to downsample the feature maps generated by the convolutional layers. They reduce the spatial dimensions of the feature maps while retaining the most salient information. Common pooling operations include but are not limited to max pooling and average pooling. Pooling layers help in achieving translation invariance, reducing computational complexity, and controlling the receptive field of the autoencoder. Rectified Linear Unit (ReLU) functions introduce non-linearity into the autoencoder by applying a ReLU activation function element-wise to the output of the previous layer. ReLU functions help in capturing complex patterns and relationships in the data by allowing the network to learn non-linear transformations. They also promote sparsity and alleviate the vanishing gradient problem. The bottleneck layer represents the most compressed representation of the input data. The bottleneck layer has a significantly reduced dimensionality compared to the input and output layers of the autoencoder. It forces the network to learn a compact and meaningful encoding of the data, capturing the essential features and discarding redundant information. In one embodiment, the multi-layer autoencoder network is comprised of a plurality of the previously mentioned layers where the sequence and composition of the layers may vary depending on a user's preferences and goals. The bottleneck layer is where the compressed output **140** is created. Each layer previous to the bottleneck layer creates a more and more compressed version of the original input. The layers after the bottleneck layer represent the decoder network **132** where a plurality of layers operate on a compressed input to decompress a data set. Decompression results in a version of the original input which is largely similar but has some lost data from the transformations.

[0055] FIG. 3 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a correlation network. The correlation network **160** is designed to enhance the reconstruction of decompressed data by leveraging correlations and patterns within the data. The correlation network **160** may also be referred to as a neural upsampler. The correlation network **160** comprises a plurality of correlation network elements that work together to capture and utilize the correlations for improved data reconstruction. Each correlation network element within the correlation network **160** contributes to the correlation learning and data reconstruction process. These elements include, but are not limited to, convolutional layers, skip connections, pooling layers and activation functions such as but not limited to, rectified linear unit functions or sigmoid functions.

[0056] The convolutional layers are responsible for extracting meaningful features from the input data. They apply convolutional operations using learnable filters to capture spatial patterns and hierarchical representations of the data. The convolutional layers can have different numbers of filters, kernel sizes, and strides to capture features at various scales and resolutions. Skip connections are employed to facilitate the flow of information across different layers of the autoencoder. Skip connections allow the output of a layer to be directly added to the output of a subsequent layer, enabling the network to learn residual mappings and mitigate the vanishing gradient problem. Skip connections help in preserving fine-grained details and improving the training stability of the autoencoder.

[0057] Pooling layers are used to downsample the feature maps generated by the convolutional layers. They reduce the spatial dimensions of the feature maps while retaining the most salient information. Common pooling operations include but are not limited to max pooling and average pooling. Pooling layers help in achieving translation invariance, reducing computational complexity, and controlling the receptive field of the autoencoder. Rectified Linear Unit (ReLU) functions introduce non-linearity into the autoencoder by applying a ReLU activation function element-wise to the output of the previous layer. ReLU functions help in capturing complex patterns and relationships in the data by allowing the network to learn non-linear transformations. They also promote sparsity and alleviate the vanishing gradient problem.

[0058] In one embodiment, the correlation network 160 may comprise an encoder 310, a decoder 320, an N number of correlated data sets 300, an N number-channel wise transformer 330, and an N number of restored data sets 340. Additionally, the correlation network 160 may be comprised of a plurality of convolutional layers, pooling layers, and activation functions. In one embodiment, the correlation network 160 may be configured to receive N correlated data sets 300 where each correlated data set includes a plurality of decompressed data points. In one embodiment, the correlation network 160 may be configured to receive four correlated data sets as an input. The correlated data sets may have been organized by a decompressed output organizer 170 to maximize the similarities between the data points in each set. One data set, 300, may include data points 300a, 300b, 300c, through 300n, where the decompressed output organizer 170 has determined the N number of data points are similar enough to be grouped together. The correlation network 160 may then receive and process full data sets at a time. In FIG. 3, the data is processed through an encoder 310 by passing through a convolutional layer, a pooling layer, and an activation function.

[0059] Activation functions introduce non-linearity into the network, enabling it to learn and represent complex patterns and relationships in the data. Common activation functions include but are not limited to sigmoid, tanh, ReLU (Rectified Linear Unit), and its variants. These functions have different properties and are chosen based on the specific requirements of the task and the network architecture. For example, ReLU is widely used in deep neural networks due to its ability to alleviate the vanishing gradient problem and promote sparsity in the activations. By applying activation functions, the neural network can learn capture non-linear relationships in the data, enabling it to model complex patterns and make accurate predictions or decisions.

[0060] The encoder 310 breaks the decompressed outputs passed by the decompressed output organizer 170 down into smaller representations of the original data sets. Following the encoder the data may pass through a transformer 330. A transformer is a type of neural network architecture that may rely on a self-attention mechanism which allows the model to weigh the importance of different parts of the input sequence when processing each element. This enables the transformer to capture dependencies and relationships between elements in the sequence efficiently. After being processed by a transformer 330, the data sets may be further processed by a decoder 320 which restores the smaller representations back into the original decompressed data sets. The decoder 320 may have a similar composition as the

encoder 310, but reversed, to undo the operations performed on the data sets by the encoder 310. The transformer 330 may identify important aspects in each group of decompressed data passed through the correlation network which allows the decoder 320 to rebuild a more complete version of the original decompressed data sets. The decoder 320 may output an N number of restored data sets 340 which correspond to the N number of correlated data sets 300 originally passed through the correlation network 170.

[0061] FIG. 4 is a block diagram illustrating an exemplary aspect of a platform for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, an autoencoder training system 270. According to the embodiment, the autoencoder training system 270 may comprise a model training stage comprising a data preprocessor 402, one or more machine and/or deep learning algorithms 403, training output 404, and a parametric optimizer 405, and a model deployment stage comprising a deployed and fully trained model 410 configured to perform tasks described herein such as transcription, summarization, agent coaching, and agent guidance. Autoencoder training system 270 may be used to train and deploy the multi-layer autoencoder network 210 in order to support the services provided by the compression and restoration system.

[0062] At the model training stage, a plurality of training data 401 may be received at the autoencoder training system 270. In some embodiments, the plurality of training data may be obtained from one or more storage systems 150 and/or directly from various information sources.

[0063] In a use case directed to hyperspectral images, a plurality of training data may be sourced from data collectors including but not limited to satellites, airborne sensors, unmanned aerial vehicles, ground-based sensors, and medical devices. Hyperspectral data refers to data that includes wide ranges of the electromagnetic spectrum. It could include information in ranges including but not limited to the visible spectrum and the infrared spectrum. Data preprocessor 402 may receive the input data (e.g., hyperspectral data) and perform various data preprocessing tasks on the input data to format the data for further processing. For example, data preprocessing can include, but is not limited to, tasks related to data cleansing, data deduplication, data normalization, data transformation, handling missing values, feature extraction and selection, mismatch handling, and/or the like. Data preprocessor 402 may also be configured to create training dataset, a validation dataset, and a test set from the plurality of input data 401. For example, a training dataset may comprise 80% of the preprocessed input data, the validation set 10%, and the test dataset may comprise the remaining 10% of the data. The preprocessed training dataset may be fed as input into one or more machine and/or deep learning algorithms 403 to train a predictive model for object monitoring and detection.

[0064] During model training, training output 404 is produced and used to measure the quality and efficiency of the compressed outputs. During this process a parametric optimizer 405 may be used to perform algorithmic tuning between model training iterations. Model parameters and hyperparameters can include, but are not limited to, bias, train-test split ratio, learning rate in optimization algorithms (e.g., gradient descent), choice of optimization algorithm (e.g., gradient descent, stochastic gradient descent, of Adam optimizer, etc.), choice of activation function in a neural

network layer (e.g., Sigmoid, ReLu, Tanh, etc.), the choice of cost or loss function the model will use, number of hidden layers in a neural network, number of activation unites in each layer, the drop-out rate in a neural network, number of iterations (epochs) in a training the model, number of clusters in a clustering task, kernel or filter size in convolutional layers, pooling size, batch size, the coefficients (or weights) of linear or logistic regression models, cluster centroids, and/or the like. Parameters and hyperparameters may be tuned and then applied to the next round of model training. In this way, the training stage provides a machine learning training loop.

[0065] In some implementations, various accuracy metrics may be used by the autoencoder training system 270 to evaluate a model's performance. Metrics can include, but are not limited to, compression ratio, the amount of data lost, the size of the compressed file, and the speed at which data is compressed, to name a few. In one embodiment, the system may utilize a loss function 407 to measure the system's performance. The loss function 407 compares the training outputs with an expected output and determined how the algorithm needs to be changed in order to improve the quality of the model output. During the training stage, all outputs may be passed through the loss function 407 on a continuous loop until the algorithms 403 are in a position where they can effectively be incorporated into a deployed model 415.

[0066] The test dataset can be used to test the accuracy of the model outputs. If the training model is compressing or decompressing data to the user's preferred standards, then it can be moved to the model deployment stage as a fully trained and deployed model 410 in a production environment compressing or decompressing live input data 411 (e.g., hyperspectral data). Further, model compressions or decompressions made by deployed model can be used as feedback and applied to model training in the training stage, wherein the model is continuously learning over time using both training data and live data and predictions.

[0067] A model and training database 406 is present and configured to store training/test datasets and developed models. Database 406 may also store previous versions of models. According to some embodiments, the one or more machine and/or deep learning models may comprise any suitable algorithm known to those with skill in the art including, but not limited to: LLMs, generative transformers, transformers, supervised learning algorithms such as: regression (e.g., linear, polynomial, logistic, etc.), decision tree, random forest, k-nearest neighbor, support vector machines, Naïve-Bayes algorithm; unsupervised learning algorithms such as clustering algorithms, hidden Markov models, singular value decomposition, and/or the like. Alternatively, or additionally, algorithms 403 may comprise a deep learning algorithm such as neural networks (e.g., recurrent, convolutional, long short-term memory networks, etc.). In some implementations, the autoencoder training system 270 automatically generates standardized model scorecards for each model produced to provide rapid insights into the model and training data, maintain model provenance, and track performance over time. These model scorecards provide insights into model framework(s) used, training data, training data specifications such as chip size, stride, data splits, baseline hyperparameters, and other factors. Model scorecards may be stored in database(s) 406.

[0068] FIG. 5 is a block diagram illustrating an exemplary aspect of a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a correlation network training system 370. According to the embodiment, correlation network training system 370 may comprise a model training stage comprising a data preprocessor 502, one or more machine and/or deep learning algorithms 503, training output 504, and a parametric optimizer 505, and a model deployment stage comprising a deployed and fully trained model 510 configured to perform tasks described herein such determining correlations between compressed data sets. The correlation network training system 370 may be used to train and deploy the correlation network 300 in order to support the services provided by the compression and decompression system.

[0069] At the model training stage, a plurality of training data 501 may be received by the correlation network training system 500. In some embodiments, the plurality of training data may be obtained from one or more storage systems 150 and/or directly from the compression network 131. In some embodiments, the correlation network training system may obtain data sets from a vector grouping system 180. In a use case directed to hyperspectral data sets, a plurality of decompressed training data may be sourced from a hyperspectral data compression system. Data preprocessor 502 may receive the input data (e.g., decompressed hyperspectral data) and perform various data preprocessing tasks on the input data to format the data for further processing. For example, data preprocessing can include, but is not limited to, tasks related to data cleansing, data deduplication, data normalization, data transformation, handling missing values, feature extraction and selection, mismatch handling, and/or the like. Data preprocessor 502 may also be configured to create training dataset, a validation dataset, and a test set from the plurality of input data 501. For example, a training dataset may comprise 80% of the preprocessed input data, the validation set 10%, and the test dataset may comprise the remaining 10% of the data. The preprocessed training dataset may be fed as input into one or more machine and/or deep learning algorithms 503 to train a predictive model for object monitoring and detection.

[0070] During model training, training output 504 is produced and used to measure the accuracy and usefulness of the predictive outputs. During this process a parametric optimizer 505 may be used to perform algorithmic tuning between model training iterations. Model parameters and hyperparameters can include, but are not limited to, bias, train-test split ratio, learning rate in optimization algorithms (e.g., gradient descent), choice of optimization algorithm (e.g., gradient descent, stochastic gradient descent, of Adam optimizer, etc.), choice of activation function in a neural network layer (e.g., Sigmoid, ReLu, Tanh, etc.), the choice of cost or loss function the model will use, number of hidden layers in a neural network, number of activation unites in each layer, the drop-out rate in a neural network, number of iterations (epochs) in a training the model, number of clusters in a clustering task, kernel or filter size in convolutional layers, pooling size, batch size, the coefficients (or weights) of linear or logistic regression models, cluster centroids, and/or the like. Parameters and hyperparameters may be tuned and then applied to the next round of model training. In this way, the training stage provides a machine learning training loop.

[0071] In some implementations, various accuracy metrics may be used by machine learning engine 400 to evaluate a model's performance. Metrics can include, but are not limited to, word error rate (WER), word information loss, speaker identification accuracy (e.g., single stream with multiple speakers), inverse text normalization and normalization error rate, punctuation accuracy, timestamp accuracy, latency, resource consumption, custom vocabulary, sentence-level sentiment analysis, multiple languages supported, cost-to-performance tradeoff, and personal identifying information/payment card industry redaction, to name a few. In one embodiment, the system may utilize a loss function 507 to measure the system's performance. The loss function 507 compares the training outputs with an expected output and determined how the algorithm needs to be changed in order to improve the quality of the model output. During the training stage, all outputs may be passed through the loss function 507 on a continuous loop until the algorithms 503 are in a position where they can effectively be incorporated into a deployed model 515.

[0072] The test dataset can be used to test the accuracy of the model outputs. If the training model is establishing correlations that satisfy a certain criterion such as but not limited to quality of the correlations and amount of restored lost data, then it can be moved to the model deployment stage as a fully trained and deployed model 510 in a production environment making predictions based on live input data 511 (e.g., compressed hyperspectral data). Further, model correlations and restorations made by deployed model can be used as feedback and applied to model training in the training stage, wherein the model is continuously learning over time using both training data and live data and predictions. A model and training database 506 is present and configured to store training/test datasets and developed models. Database 506 may also store previous versions of models.

[0073] According to some embodiments, the one or more machine and/or deep learning models may comprise any suitable algorithm known to those with skill in the art including, but not limited to: LLMs, generative transformers, transformers, supervised learning algorithms such as: regression (e.g., linear, polynomial, logistic, etc.), decision tree, random forest, k-nearest neighbor, support vector machines, Naïve-Bayes algorithm; unsupervised learning algorithms such as clustering algorithms, hidden Markov models, singular value decomposition, and/or the like. Alternatively, or additionally, algorithms 503 may comprise a deep learning algorithm such as neural networks (e.g., recurrent, convolutional, long short-term memory networks, etc.).

[0074] In some implementations, the correlation network training system 270 automatically generates standardized model scorecards for each model produced to provide rapid insights into the model and training data, maintain model provenance, and track performance over time. These model scorecards provide insights into model framework(s) used, training data, training data specifications such as chip size, stride, data splits, baseline hyperparameters, and other factors. Model scorecards may be stored in database(s) 506.

Detailed Description of Exemplary Aspects

[0075] FIG. 6 is a flow diagram illustrating an exemplary method for compressing a data input using a system for compressing and restoring data using multi-level autoencod-

ers and correlation networks. In a first step 600, a plurality of data sets is collected from a plurality of data sources. These data sources can include various sensors, devices, databases, or any other systems that generate or store data. The data sets may be heterogeneous in nature, meaning they can have different formats, structures, or modalities. For example, the data sets can include images, videos, audio recordings, time-series data, numerical data, or textual data. The collection process involves acquiring the data sets from their respective sources and bringing them into a centralized system for further processing.

[0076] In a step 610, the collected data sets are preprocessed using a data preprocessor. The data preprocessor may be responsible for cleaning, transforming, and preparing the data sets for subsequent analysis and compression. Preprocessing tasks may include but are not limited to data cleansing, data integration, data transformation, and feature extraction. Data cleansing involves removing or correcting any erroneous, missing, or inconsistent data points. Data integration combines data from multiple sources into a unified format. Data transformation converts the data into a suitable representation for further processing, such as scaling, normalization, or encoding categorical variables. Feature extraction identifies and selects relevant features or attributes from the data sets that are most informative for the given task.

[0077] A step 620 involves normalizing the preprocessed data sets using a data normalizer. Normalization is a step that brings the data into a common scale and range. It helps to remove any biases or inconsistencies that may exist due to different units or scales of measurement. The data normalizer applies various normalization techniques, such as min-max scaling, z-score normalization, or unit vector normalization, depending on the nature of the data and the requirements of the subsequent compression step. Normalization ensures that all the data sets have a consistent representation and can be compared and processed effectively.

[0078] In a step 630, the normalized data sets are compressed into a compressed output using a multi-layer auto-encoder network. The multi-layer autoencoder network is a deep learning model designed to learn compact and meaningful representations of the input data. It consists of an encoder network and a decoder network. The encoder network takes the normalized data sets as input and progressively compresses them through a series of layers, such as but not limited to convolutional layers, pooling layers, and fully connected layers. The compressed representation is obtained at the bottleneck layer of the encoder network, which has a significantly reduced dimensionality compared to the original data. The multi-layer autoencoder network may utilize a plurality of encoder networks to achieve optimal compression performance. These encoder networks can include different architectures, loss functions, or optimization techniques. The choice of compression technique depends on the specific characteristics and requirements of the data sets being compressed. During the compression process, the multi-layer autoencoder network learns to capture the essential features and patterns present in the data sets while discarding redundant or irrelevant information. It aims to minimize the reconstruction error between the original data and the reconstructed data obtained from the compressed representation. In step 640, the compressed output generated by the multi-layer autoencoder network is

either outputted or stored for future processing. The compressed output represents the compact and informative representation of the original data sets. It can be transmitted, stored, or further analyzed depending on the specific application or use case. The compressed output significantly reduces the storage and transmission requirements compared to the original data sets, making it more efficient for downstream tasks.

[0079] FIG. 7 is a flow diagram illustrating an exemplary method for decompressing a compressed data input using system for compressing and restoring data using multi-level autoencoders and correlation networks. In a first step, **700**, access a plurality of compressed data sets. In a step **710**, decompress the plurality of compressed data sets using a multi-layer autoencoder's decoder network. The decoder network is responsible for mapping the latent space vectors back to the original data space. The decoder network may include techniques such as transposed convolutions, upsampling layers, or generative models, depending on the specific requirements of the data and the compression method used.

[0080] In a step **720**, leverage the similarities between decompressed outputs using a correlation network which may exploit shared information and patterns to achieve a better reconstruction. The correlation network is a deep learning model specifically designed to exploit the shared information and patterns among the compressed data sets. It takes the organized decompressed data sets as input and learns to capture the correlations and dependencies between them. The correlation network may consist of multiple layers, such as convolutional layers, recurrent layers, or attention mechanisms, which enable it to effectively model the relationships and similarities among the compressed data sets.

[0081] In a step **730**, the compressed data sets are reconstructed using the correlation network. The reconstruction process in step **730** combines the capabilities of the correlation network and the decompression systems. The correlation network provides the enhanced and refined latent space representations, while the decompression systems use these representations to generate the reconstructed data. In a step **740**, the restored, decompressed data set is outputted. The restored data set represents the reconstructed version of the original data, which includes recovered information lost during the compression process. The outputted data set more closely resembles the original data than would a decompressed output passed solely through a decoder network.

[0082] FIG. 8 is a block diagram illustrating an exemplary system architecture for compressing and restoring IoT sensor data using a system for compressing and restoring data using multi-level autoencoders and correlation networks. The IoT Sensor Stream Organizer **800** is responsible for collecting and organizing data streams from various IoT sensors. It receives raw sensor data from multiple sources, such as but not limited to temperature sensors, humidity sensors, and accelerometers. The IoT Sensor Stream Organizer **800** may perform necessary preprocessing tasks, such as data cleaning, normalization, and synchronization, to ensure the data is in a suitable format for further processing. The preprocessed IoT sensor data is then passed to a data preprocessor **810**. The data preprocessor **810** prepares the data for compression by transforming it into a latent space representation. It applies techniques such as feature extraction, dimensionality reduction, and data normalization to extract meaningful features and reduce the dimensionality of

the data. The latent space representation captures the essential characteristics of the IoT sensor data while reducing its size.

[0083] The multi-layer autoencoder **820** is responsible for compressing and decompressing the latent space representation of the IoT sensor data. It consists of an encoder network **821** and a decoder network **822**. The encoder network **821** takes the latent space representation as input and progressively compresses it through a series of layers, such as but not limited to convolutional layers, pooling layers, and fully connected layers. The compressed representation may pass through a bottleneck layer which transforms the original data to have a significantly reduced dimensionality compared to the original data. Further, the encoder network **821** manages the compression process and stores the compressed representation of the IoT sensor data. It determines the optimal compression settings based on factors such as the desired compression ratio, data characteristics, and available storage resources. The compressed representation is efficiently stored or transmitted, reducing the storage and bandwidth requirements for IoT sensor data.

[0084] The decoder network **822** is responsible for reconstructing the original IoT sensor data from the compressed representation. It utilizes the multi-layer autoencoder **820** to map the compressed representation back to the original data space. The decoder network consists of layers such as transposed convolutional layers, upsampling layers, and fully connected layers. It learns to reconstruct the original data by minimizing the reconstruction error between the decompressed output and the original IoT sensor data. The decompressed output **850** represents the decompressed IoT sensor data obtained from the decoder network **822**. It closely resembles the original data and retains the essential information captured by the sensors, but includes some information lost during the compression process. The decompressed output **850** may be further processed, analyzed, or utilized by downstream applications or systems.

[0085] To further enhance the compression and reconstruction quality, the system includes a correlation network **830**. The correlation network **830** learns and exploits correlations and patterns within the IoT sensor data to improve the reconstruction process. It consists of multiple correlation layers that capture dependencies and relationships among different sensors or data streams. The correlation network **830** helps in preserving important information that may have been lost during the compression process. Following the identification of dependencies and relationships among different data streams, the correlation network **830** reconstruct a decompressed output **850** into a restored output **860** which recovers much of the data lost during the compression and decompression process.

[0086] The system may be trained using an end-to-end approach, where the multi-layer autoencoder **820** and the correlation network **830** are jointly optimized to minimize the reconstruction error and maximize the compression ratio. The training process may involve feeding the IoT sensor data through the system, comparing the decompressed output with the original data, and updating the network parameters using backpropagation and gradient descent techniques. The proposed system offers several advantages for IoT sensor data compression. It achieves high compression ratios while preserving the essential information in the data. The multi-layer autoencoder **820** learns compact and meaningful representations of the data, exploit-

ing spatial and temporal correlations. The correlation network **830** further enhances the compression quality by capturing dependencies and patterns within the data. Moreover, the system is adaptable and can handle various types of IoT sensor data, making it suitable for a wide range of IoT applications. It can be deployed on resource-constrained IoT devices or edge servers, reducing storage and transmission costs while maintaining data quality.

[0087] FIG. 9 is a flow diagram illustrating an exemplary method for compressing and decompressing IoT sensor data using a system for compressing and restoring data using multi-level autoencoders and correlation networks. In a first step **900**, incoming IoT sensor data is organized based on its origin sensor type. IoT sensor data can be generated from various types of sensors, such as but not limited to temperature sensors, humidity sensors, pressure sensors, accelerometers, or any other sensors deployed in an IoT network. Each sensor type captures specific measurements or data points relevant to its function. The organization step involves categorizing and grouping the incoming IoT sensor data based on the type of sensor it originated from. This step helps to maintain a structured and organized representation of the data, facilitating subsequent processing and analysis.

[0088] In a step **910**, the latent space vectors for each IoT sensor data set are preprocessed. Latent space vectors are lower-dimensional representations of the original data that capture the essential features and patterns. Preprocessing the latent space vectors involves applying various techniques to ensure data quality, consistency, and compatibility. This may include but is not limited to data cleaning, normalization, feature scaling, or dimensionality reduction. The preprocessing step aims to remove any noise, outliers, or inconsistencies in the latent space vectors and prepare them for the compression process.

[0089] A step **920** involves compressing each IoT sensor data set using a multi-layer autoencoder network. The multi-layer autoencoder network is a deep learning model designed to learn compact and meaningful representations of the input data. It consists of an encoder network and a decoder network. The encoder network takes the preprocessed latent space vectors as input and progressively compresses them through a series of layers, such as convolutional layers, pooling layers, and fully connected layers. The compressed representation is obtained at the bottleneck layer of the encoder network, which has a significantly reduced dimensionality compared to the original data. The multi-layer autoencoder network may include a compression system that specifically handles the compression of IoT sensor data. The compression system can employ various techniques, such as quantization, entropy coding, or sparse representations, to achieve efficient compression while preserving the essential information in the data. The compression system outputs a compressed IoT sensor data set, which is a compact representation of the original data. In step **930**, the original IoT sensor data is decompressed using a decoder network. The decoder network is responsible for reconstructing the original data from the compressed representation. It takes the compressed IoT sensor data sets and applies a series of decompression operations, such as transposed convolutions or upsampling layers, to map the compressed data back to its original dimensionality.

[0090] In a step **940**, correlations between compressed IoT sensor data sets are identified using a correlation network. The correlation network is a separate deep learning model

that learns to capture the relationships and dependencies among different compressed IoT sensor data sets. It takes the decompressed data sets as input and identifies patterns, similarities, and correlations among them. The correlation network can utilize techniques such as convolutional layers, attention mechanisms, or graph neural networks to effectively model the interactions and dependencies between the compressed data sets. The identified correlations provide valuable insights into how different IoT sensor data sets are related and how they influence each other. These correlations can be used to improve the compression efficiency and enhance the restoration quality of the data.

[0091] In a step **950**, the correlation network creates a restored, more reconstructed version of the decompressed output. By leveraging correlations between decompressed outputs, the correlation network is able to recover a large portion of information lost during the compression and decompression process. The restored, reconstructed output is similar to the decompressed output and the original input, but recovers information that may have been missing in the decompressed output.

[0092] FIG. 10 is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, the decompressed output organizer. In one embodiment, the decompressed output organizer **170** may create a matrix of n by n data sets where each data set represents a decompressed set of information. In the embodiment depicted, the decompressed output organizer **170** outputs a 4 by 4 matrix of decompressed data sets. The organizer **170** may organize the decompressed data sets into groups based on how correlated each data set is to each other. For example, decompressed data set 1 which includes **1000a**, **1000b**, **1000c**, and **1000n**, is a set of four data sets that the decompressed output organizer **170** has determined to be highly correlated. The same is true for decompressed data sets 2, 3, and 4.

[0093] The decompressed output organizer primes the correlation network **160** to receive an already organizer plurality of inputs. The correlation network may take a plurality of decompressed data sets as its input, depending on the size of the organized matrix produced by the decompressed output organizer **170**. For example, in the embodiment depicted in FIG. 10, the decompressed output organizer **170** produces a 4 by 4 matrix of data sets. The correlation network in turn receives a 4 element data set as its input. If decompressed data set 1 were to be processed by the correlation network **160**, the correlation network **160** may take **1000a**, **1000b**, **1000c**, and **1000n**, as the inputs and process all four data sets together. By clustering data sets together into groups based on how correlated they are, the decompressed output organizer **170** allows the correlation network **160** to produce more outputs that better encompass the original pre-compressed and decompressed data sets. More information may be recovered by the correlation network **160** when the inputs are already highly correlated.

[0094] FIG. 11 is a flow diagram illustrating an exemplary method for organizing restored, decompressed data sets after correlation network processing. In a first step **1100**, access a plurality of restored data sets. In a step **1110**, organize the plurality of restored data sets based on similarities if necessary. In a step **1120**, output a plurality of restored, potentially organizer data sets. This method essentially reassesses the organizational grouping performed by the decompressed

output organizer **170**. The correlation network **160** may output a matrix where the matrix contains a plurality of restored, decompressed data sets. The final output of the system may reorganize the restored, decompressed data sets within the outputted matrix based on user preference and the correlations between each data set within the matrix.

[0095] FIG. **12** is a block diagram illustrating an exemplary system architecture for compressing and restoring data using hierarchical autoencoders and correlation networks. This network replaces the previous single-level autoencoder, offering improved compression and decompression capabilities across multiple scales of data features.

[0096] A hierarchical autoencoder network **1200** comprises of two main components: a hierarchical encoder network **1210** and a hierarchical decoder network **1220**. The hierarchical encoder network **1210** comprises multiple levels of encoders, each designed to capture and compress features at different scales. As data flows through the encoder levels, it is progressively compressed, with each level focusing on increasingly fine-grained features of the input data.

[0097] When compressing data, the system first processes the input data **100** through the data preprocessor **110** and data normalizer **120**. The normalized data then enters the hierarchical encoder network **1210**. The first level of the encoder captures large-scale features, passing its output to the second level, which focuses on medium-scale features. This process continues through subsequent levels, each concentrating on finer details. The final output of the hierarchical encoder network is a multi-level compressed representation, stored as the compressed output **140**.

[0098] For decompression, the system utilizes the hierarchical decoder network **1220**. This network mirrors the structure of the encoder but operates in reverse. The compressed output **140** enters the highest level of the decoder, which begins reconstructing the coarsest features. Each subsequent level of the decoder adds finer details to the reconstruction, using both the output from the previous level and the corresponding level's compressed representation. The final level of the decoder produces the decompressed output **170**.

[0099] In one embodiment, hierarchical encoder network **1210** may include of several levels, each designed to capture features at different scales. For instance, in a four-level encoder, Level 1 (largest scale) might use convolutional layers with large kernels and aggressive pooling to capture the most general features of the image, such as overall color distribution and major structural elements. It could reduce the image to, say, $\frac{1}{8}$ of its original dimensions. A level 2 (medium scale) which operates on the output of Level 1 might use slightly smaller kernels and less aggressive pooling to capture medium-scale features like edges and basic shapes. It might further reduce the representation to $\frac{1}{4}$ of Level 1's output. A level 3 (fine scale) could focus on more detailed features, potentially using dilated convolutions to capture longer-range dependencies without further dimension reduction. A level 4 (finest scale) might use very small kernels to capture the finest details and textures in the image, with minimal or no further dimension reduction. The compressed output **140** would be a combination of the outputs from all these levels, providing a multi-scale representation of the original image.

[0100] Similarly, hierarchical decoder network **1220** would mirror this structure in reverse. A level 4 decoder may

start with the finest-scale compressed representation that begins reconstructing the detailed features and textures. A level 3 decoder may combine its input with the Level 3 compressed representation, adding finer details to the reconstruction. A level 2 decoder may utilize upsampling or transposed convolutions. This level begins to restore the spatial dimensions, adding medium-scale features back into the image. A final level further upsamples the image, restoring it to its original dimensions and reconstructing the coarsest features. Each decoder level would combine the output from the previous level with the corresponding encoder level's output, allowing for the progressive restoration of details at each scale.

[0101] This multi-level approach allows the system to efficiently compress and accurately reconstruct features at various scales, potentially leading to better overall compression performance and reconstruction quality compared to a single-level autoencoder. The system integrates a hierarchical autoencoder trainer **1230** to optimize the performance of the hierarchical autoencoder network **1200**. This trainer adjusts the parameters of both the encoder and decoder networks across all levels, ensuring efficient compression and accurate reconstruction for various types of input data. After decompression, the system continues to employ the decompressed output organizer **190** and correlation network **160** leveraging multi-scale correlations to further enhance the reconstructed output **180**. This hierarchical approach allows the system to adapt to a wide range of data types and scales, potentially achieving higher compression ratios while maintaining or improving the quality of data restoration.

[0102] FIG. **13** is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using hierarchical autoencoders and correlation networks, a hierarchical autoencoder. This multi-level architecture enables the system to process and represent data at various scales, leading to more efficient compression and higher-quality restoration.

[0103] The hierarchical encoder network **1210** comprises multiple encoding levels. In the illustrated example the hierarchical encoder network includes 3 layers of both encoders and decoders to effectively compress and decompress various levels of input representations. Each level is designed to capture and compress different aspects of the input data:

[0104] A level 1 encoder **1300** focuses on large-scale, global features of the input. For instance, in image data, this level might capture overall color distributions, major structural elements, or low-frequency patterns. It employs techniques such as large convolutional kernels or aggressive pooling to downsample the input significantly. A level 2 encoder **1310** processes the output from level 1 encoder **1300**, concentrating on medium-scale features. It might identify edges, basic shapes, or texture patterns. This level typically uses smaller convolutional kernels and less aggressive pooling, striking a balance between feature detail and data reduction.

[0105] A level 3 encoder **1320** targets fine-grained details in the data. It might employ techniques like dilated convolutions to capture intricate patterns or long-range dependencies without further downsampling. This level preserves the highest frequency components of the input that are still relevant for reconstruction. The outputs from all encoder levels are combined to form the compressed output **140**,

which is a multi-scale representation of the original input. This approach allows the system to retain important information at various scales, facilitating more effective compression than a single-scale approach.

[0106] The hierarchical decoder network **1220** mirrors the encoder structure, with level 3 decoder **1330**, level 2 decoder **1340**, and level 1 decoder **1350**. These decoders work in concert to progressively reconstruct the original input. Level 3 decoder **1330** begins the reconstruction process using the finest-scale information from the compressed output. It may employ techniques like small convolutional kernels or attention mechanisms to start rebuilding detailed features. A level 2 decoder **1340** combines its input with corresponding information from the compressed output to add medium-scale features. It may use upsampling or transposed convolutions to increase spatial dimensions, reconstructing shapes and edges.

[0107] Level 1 decoder **1350** finalizes the reconstruction, integrating coarse-scale features and restoring the output to its original dimensions. It may use larger convolutional kernels or sophisticated upsampling techniques to ensure smooth, globally coherent outputs. In one embodiment, each decoder level not only uses the output from the previous decoder level but also incorporates the corresponding encoder level's output. This creates short-cut connections that allow high-fidelity details to flow more directly from input to output, potentially improving reconstruction quality.

[0108] By leveraging different aspects of the input at each level, this hierarchical structure allows the system to compress data more efficiently and restore it more accurately. Coarse levels capture global structure and context, while finer levels preserve intricate details. This multi-scale approach enables the system to adapt to various types of data and achieve a better balance between compression ratio and reconstruction quality compared to single-scale methods. The decompressed output **170** produced by this hierarchical process retains both large-scale structures and fine details of the original input, providing a high-quality reconstruction that can be further refined by subsequent stages of the system, such as the correlation network.

[0109] FIG. 14 is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using hierarchical autoencoders and correlation networks, a hierarchical autoencoder trainer. According to the embodiment, the hierarchical autoencoder training system **1230** may comprise a model training stage comprising a data preprocessor **1402**, one or more machine and/or deep learning algorithms **1403**, training output **1404**, and a parametric optimizer **1405**, and a model deployment stage comprising a deployed and fully trained model **1410** configured to perform tasks described herein such as transcription, summarization, agent coaching, and agent guidance. Hierarchical autoencoder trainer **1230** may be used to train and deploy the hierarchical autoencoder network **1200** in order to support the services provided by the compression and restoration system.

[0110] At the model training stage, a plurality of training data **401** may be received at the hierarchical autoencoder trainer **270**. In a use case directed to hyperspectral images, a plurality of training data may be sourced from data collectors including but not limited to satellites, airborne sensors, unmanned aerial vehicles, ground-based sensors, and medical devices. Hyperspectral data refers to data that includes wide ranges of the electromagnetic spectrum. It

could include information in ranges including but not limited to the visible spectrum and the infrared spectrum. Data preprocessor **1402** may receive the input data (e.g., hyperspectral data, text data, image data, audio data) and perform various data preprocessing tasks on the input data to format the data for further processing. For example, data preprocessing can include, but is not limited to, tasks related to data cleansing, data deduplication, data normalization, data transformation, handling missing values, feature extraction and selection, mismatch handling, and/or the like. Data preprocessor **1402** may also be configured to create training dataset, a validation dataset, and a test set from the plurality of input data **401**. For example, a training dataset may comprise 80% of the preprocessed input data, the validation set 10%, and the test dataset may comprise the remaining 10% of the data. The preprocessed training dataset may be fed as input into one or more machine and/or deep learning algorithms **1403** to train a predictive model for object monitoring and detection.

[0111] During model training, training output **1404** is produced and used to measure the quality and efficiency of the compressed outputs. During this process a parametric optimizer **1405** may be used to perform algorithmic tuning between model training iterations. Model parameters and hyperparameters can include, but are not limited to, bias, train-test split ratio, learning rate in optimization algorithms (e.g., gradient descent), choice of optimization algorithm (e.g., gradient descent, stochastic gradient descent, of Adam optimizer, etc.), choice of activation function in a neural network layer (e.g., Sigmoid, ReLu, Tanh, etc.), the choice of cost or loss function the model will use, number of hidden layers in a neural network, number of activation unites in each layer, the drop-out rate in a neural network, number of iterations (epochs) in a training the model, number of clusters in a clustering task, kernel or filter size in convolutional layers, pooling size, batch size, the coefficients (or weights) of linear or logistic regression models, cluster centroids, and/or the like. Parameters and hyperparameters may be tuned and then applied to the next round of model training. In this way, the training stage provides a machine learning training loop.

[0112] In some implementations, various accuracy metrics may be used by the hierarchical autoencoder trainer **1230** to evaluate a model's performance. Metrics can include, but are not limited to, compression ratio, the amount of data lost, the size of the compressed file, and the speed at which data is compressed, to name a few. In one embodiment, the system may utilize a loss function **407** to measure the system's performance. The loss function **407** compares the training outputs with an expected output and determined how the algorithm needs to be changed in order to improve the quality of the model output. During the training stage, all outputs may be passed through the loss function **407** on a continuous loop until the algorithms **1403** are in a position where they can effectively be incorporated into a deployed model **1415**.

[0113] The test dataset can be used to test the accuracy of the model outputs. If the training model is compressing or decompressing data to the user's preferred standards, then it can be moved to the model deployment stage as a fully trained and deployed model **1410** in a production environment compressing or decompressing live input data **1411** (e.g., hyperspectral data, text data, image data, video data, audio data). Further, model compressions or decompressions

made by deployed model can be used as feedback and applied to model training in the training stage, wherein the model is continuously learning over time using both training data and live data and predictions.

[0114] A model and training database **1406** is present and configured to store training/test datasets and developed models. Database **1406** may also store previous versions of models.

[0115] According to some embodiments, the one or more machine and/or deep learning models may comprise any suitable algorithm known to those with skill in the art including, but not limited to: LLMs, generative transformers, transformers, supervised learning algorithms such as: regression (e.g., linear, polynomial, logistic, etc.), decision tree, random forest, k-nearest neighbor, support vector machines, Naïve-Bayes algorithm; unsupervised learning algorithms such as clustering algorithms, hidden Markov models, singular value decomposition, and/or the like. Alternatively, or additionally, algorithms **1403** may comprise a deep learning algorithm such as neural networks (e.g., recurrent, convolutional, long short-term memory networks, etc.). In some implementations, the hierarchical autoencoder trainer **1230** automatically generates standardized model scorecards for each model produced to provide rapid insights into the model and training data, maintain model provenance, and track performance over time. These model scorecards provide insights into model framework(s) used, training data, training data specifications such as chip size, stride, data splits, baseline hyperparameters, and other factors. Model scorecards may be stored in database(s) **1406**.

[0116] In one embodiment, hierarchical autoencoder trainer **1230** may employ a flexible approach to optimize the performance of both the hierarchical encoder and decoder networks. It can train each level of the encoders and decoders either separately or jointly, adapting to the specific requirements of the system and the characteristics of the input data. When training levels separately, the trainer focuses on optimizing each level's ability to capture or reconstruct features at its particular scale, allowing for fine-tuned performance at each stage of the compression and decompression process. This approach can be particularly useful when dealing with diverse data types or when specific levels require specialized attention. Alternatively, joint training of multiple or all levels enables the system to learn inter-level dependencies and optimize the overall compression-decompression pipeline as a cohesive unit. This can lead to improved global performance and more efficient use of the network's capacity. The trainer may also employ a hybrid approach, starting with separate level training to establish baseline performance, followed by joint fine-tuning to enhance overall system coherence. This adaptive training strategy ensures that the hierarchical autoencoder network can be optimized for a wide range of applications and data types, maximizing both compression efficiency and reconstruction quality.

[0117] FIG. **15** is a flow diagram illustrating an exemplary method for compressing and restoring data using hierarchical autoencoders and correlation networks. In a first step **1500**, the system collects and preprocesses a plurality of data sets from various data sources. This initial stage involves gathering diverse data types, which may include images, videos, sensor readings, or other forms of structured or unstructured data. The preprocessing phase is crucial for preparing the data for efficient compression. It may involve

tasks such as noise reduction, normalization, and feature extraction. By standardizing the input data, this step ensures that the subsequent compression process can operate effectively across different data modalities and scales.

[0118] In a step **1510**, the system compresses the normalized data sets using a hierarchical multi-layer autoencoder network. This step marks the beginning of the advanced compression process, utilizing the sophisticated hierarchical structure of the autoencoder. The hierarchical approach allows for a more nuanced and efficient compression compared to traditional single-layer methods, as it can capture and preserve information at multiple scales simultaneously.

[0119] In a step **1520**, the system processes the data through the level 1 encoder, focusing on large-scale features, and continues through subsequent levels, each focusing on finer-scale features. This multi-level encoding is at the heart of the hierarchical compression process. The level 1 encoder captures broad, global features of the data, while each subsequent level concentrates on increasingly fine-grained details. This approach ensures that the compressed representation retains a rich, multi-scale characterization of the original data, potentially leading to better compression ratios and more accurate reconstruction.

[0120] In a step **1530**, the system outputs the multi-level compressed representation or stores it for future processing. This step represents the culmination of the compression process, where the hierarchically compressed data is either immediately utilized or securely stored. The multi-level nature of this compressed representation allows for flexible use in various applications, potentially enabling progressive decompression or scale-specific analysis without full decompression.

[0121] In a step **1540**, the system inputs the multi-level compressed representation into the hierarchical decoder when decompression is required. This step initiates the reconstruction process, leveraging the multi-scale information captured during compression. The hierarchical decoder mirrors the structure of the encoder, progressively rebuilding the data from the coarsest to the finest scales.

[0122] In a step **1550**, the system processes the decompressed output through a correlation network to restore data potentially lost during compression. This advanced restoration step goes beyond simple decompression, utilizing learned correlations at multiple scales to infer and recreate details that may have been diminished or lost in the compression process. The approach of the correlation network complements the hierarchical nature of the autoencoder, potentially leading to higher quality reconstructions.

[0123] In a step **1560**, the system outputs the restored, reconstructed data set. This final step delivers the fully processed data, which has undergone hierarchical compression, decompression, and correlation-based enhancement. The resulting output aims to closely resemble the original input data, with the potential for even enhancing certain aspects of the data through the learned correlations and multi-scale processing.

[0124] FIG. **16** is a block diagram of an exemplary system architecture for compressing and restoring data using hierarchical autoencoders and knowledge-enhanced correlation networks. The system begins with data input **100**, which represents raw data from various sources requiring compression and subsequent restoration. This data input **100** is first processed by a data preprocessor **110**, which performs various preprocessing tasks such as data cleaning, noise

reduction, and format conversion to prepare the data for subsequent processing. The preprocessed data is then passed to a data normalizer **120**, which scales and normalizes the data to a consistent range, typically between 0 and 1, to improve training stability and convergence of the autoencoder network.

[0125] The normalized data proceeds to the hierarchical autoencoder network **1200**, which comprises two primary components: a hierarchical encoder network **1210** and a hierarchical decoder network **1220**. Hierarchical encoder network **1210** is responsible for compressing the normalized data into a lower-dimensional latent space representation through multiple levels of encoding, each capturing features at different scales. This multi-level approach allows the system to efficiently represent both coarse and fine-grained features of the input data. The compressed representation from hierarchical encoder network **1210** forms compressed output **140**, which has significantly reduced dimensionality compared to the original input data.

[0126] Hierarchical decoder network **1220** performs the inverse operation, decompressing compressed output **140** to generate decompressed output **170**. Hierarchical decoder network **1220** mirrors the structure of the encoder but operates in reverse, progressively reconstructing the data from the latent representation while aiming to preserve essential characteristics of the original data. Decompressed output **170** is a reconstruction of the original data based solely on the information contained in the compressed representation.

[0127] Hierarchical autoencoder trainer **1230** optimizes the performance of hierarchical autoencoder network **1200** by adjusting the parameters of both the encoder and decoder networks across all levels. This training process ensures efficient compression and accurate reconstruction for various types of input data.

[0128] A novel aspect of the system is the incorporation of external knowledge sources **1600** into the data restoration process. External knowledge sources **1600** comprise domain-specific knowledge repositories, ontologies, rule systems, and other structured information sources that provide semantic context and constraints relevant to the data being processed. These knowledge sources enhance the system's ability to understand and preserve semantically important features during compression and restoration.

[0129] After decompression, decompressed output **170** is passed to a decompressed output organizer **190**, which groups the decompressed outputs based on their correlations and similarities. This organization step facilitates more effective restoration by knowledge graph correlation network **1610**.

[0130] Knowledge graph correlation network **1610** represents a significant advancement over the previously described correlation network. It integrates external knowledge sources **1600** with traditional correlation-based restoration techniques. Knowledge graph correlation network **1610** analyzes the decompressed data in the context of domain knowledge, identifying semantic relationships and constraints that guide the restoration process. By leveraging knowledge graphs that represent domain concepts and their relationships, the network can infer missing or degraded information more accurately than purely statistical approaches.

[0131] The output from knowledge graph correlation network **1610** is further processed by an explainable restorer

1620, which not only enhances the restoration quality but also provides transparency into the restoration process. Explainable restorer **1620** generates metadata that indicates which knowledge sources influenced specific aspects of the restoration and the confidence level associated with each restored element. This explainability aspect is particularly valuable in applications where understanding the basis for restoration decisions is critical.

[0132] The final output of the system is knowledge-enhanced reconstructed output **1630**, which represents a significantly improved reconstruction of the original data compared to initial decompressed output **170**. Knowledge-enhanced reconstructed output **1630** benefits from both the statistical patterns learned by hierarchical autoencoder network **1200** and semantic understanding provided by the external knowledge sources **1600**.

[0133] The integration of knowledge-based approaches with hierarchical autoencoding enables the system to achieve superior restoration quality across a wide range of data types and applications, particularly in domains with rich semantic structure and available knowledge resources. This combined approach addresses limitations of purely data-driven methods, especially when dealing with complex, structured data where domain knowledge provides valuable constraints and context for accurate restoration.

[0134] FIG. 17 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data, a knowledge graph correlation network. Knowledge graph correlation network **1610** receives decompressed sets **1700** as input. These decompressed sets **1700** are the output from the hierarchical decoder network and have been organized by the decompressed output organizer based on similarities and correlations. While these decompressed sets **1700** contain much of the information from the original data, they typically suffer from some information loss due to the compression process. A purpose of knowledge graph correlation network **1610** is to recover this lost information by leveraging both statistical correlations and domain knowledge.

[0135] Decompressed sets **1700** are first processed by a correlation encoder **1710**. Correlation encoder **1710** analyzes the input data to identify potential correlations and patterns within and across the decompressed sets. This component employs multiple processing stages including but not limited to feature extraction layers, correlation attention mechanisms, and latent space projection. The feature extraction layers identify characteristics and features within the decompressed data. The correlation attention mechanism focuses computational resources on areas where correlations are likely to be most informative for restoration purposes. The latent space projection transforms the extracted features into a representation that facilitates knowledge integration in subsequent stages.

[0136] In one embodiment, the system may utilize a knowledge integrator **1720**, which serves as the bridge between purely data-driven approaches and knowledge-based reasoning. Knowledge integrator **1720** receives inputs from both correlation encoder **1710** and external knowledge sources **1730**. External knowledge sources **1730** comprise domain-specific repositories, knowledge graphs, ontologies, rule systems, and other structured information resources relevant to the data domain. These might include, for

example, physical laws for scientific data, anatomical models for medical imaging, or semantic relationships for textual data.

[0137] Knowledge integrator 1720 performs several key functions. First, it maps the features and patterns identified by correlation encoder 1710 to concepts and relationships in external knowledge sources 1730. This mapping process establishes semantic meaning for the statistical patterns observed in the data. Second, it identifies relevant constraints, rules, and typical patterns from the knowledge sources that can guide the restoration process. Third, it resolves potential conflicts between statistical evidence and domain knowledge, determining how to optimally combine these different information sources. Knowledge integrator 1720 may employ techniques such as graph neural networks, attention mechanisms, and probabilistic reasoning to effectively fuse data-driven and knowledge-driven insights.

[0138] The output from knowledge integrator 1720 flows to both correlation decoder 1740 and explainable restorer 1620. Correlation decoder 1740 is responsible for transforming the knowledge-enhanced representation back into the original data space. This component includes knowledge-guided expansion layers, which leverage the integrated knowledge to infer missing or degraded information; feature reconstruction layers, which rebuild the complete feature set based on the enhanced representation; and consistency enforcement mechanisms, which ensure the restored data adheres to domain constraints and maintains internal coherence.

[0139] Explainable restorer 1620 adds a critical dimension of transparency to the restoration process. Rather than functioning as a “black box,” this component generates detailed explanations for the restoration decisions. Explainable restorer 1620 tracks which knowledge sources influenced particular aspects of the restoration, quantifies confidence levels for different restored elements, and provides a rationale for why specific restorations were made. This explainability is particularly valuable in domains where understanding the basis for restoration choices is important, such as scientific research, medical applications, or legal contexts.

[0140] Both correlation decoder 1740 and explainable restorer 1620 contribute to the final knowledge-enhanced reconstructed output 1630. This output represents a significantly improved reconstruction compared to the initial decompressed data, benefiting from both statistical correlations and domain knowledge. Knowledge-enhanced reconstructed output 1630 may also include accompanying meta-data that provides explainability and confidence metrics for the restoration process.

[0141] The architecture enables several distinctive capabilities compared to conventional restoration approaches. First, by incorporating domain knowledge, the system can restore information that would be impossible to recover using purely statistical methods. Second, the explicit knowledge integration allows the system to maintain semantic consistency and adhere to domain constraints, producing more meaningful and accurate restorations. Third, the explainability aspect provides transparency that is often lacking in deep learning-based systems, making the approach suitable for applications where understanding the restoration process is as important as the restoration itself.

[0142] Furthermore, this architecture is adaptable to various knowledge representation formats and can be special-

ized for different domains by modifying external knowledge sources 1730 without necessarily retraining the entire network. This flexibility makes the system applicable across a wide range of data types and application domains, from scientific data processing to multimedia content restoration.

[0143] FIG. 18 is a block diagram illustrating an exemplary architecture for a subsystem of the knowledge graph correlation network, a knowledge integrator. In one embodiment, knowledge integrator 1720 comprises multiple specialized components that work together to effectively incorporate external knowledge into the data restoration process. These components include but may not be limited to a knowledge graph processor 1800, a semantic mapper 1810, a rule-based constrainer 1820, a knowledge-feature aligner 1830, and a confidence estimator 1840.

[0144] Knowledge graph processor 1800 is responsible for interpreting and navigating the structured knowledge resources provided by the external knowledge sources. Knowledge graphs represent domain concepts as nodes and their relationships as edges, providing a rich, interconnected representation of domain knowledge. Knowledge graph processor 1800 performs several functions including but not limited to querying and retrieving relevant knowledge sub-graphs based on the current data context, performing reasoning operations over the knowledge graph to infer additional relationships or constraints that may not be explicitly stated, and generating embeddings or numerical representations of knowledge graph entities and relationships that can be more easily integrated with the statistical representations from the correlation encoder. Knowledge graph processor 1800 may employ techniques such as but not limited to graph neural networks, random walks, or attention mechanisms to effectively process the knowledge graph structure.

[0145] Semantic mapper 1810 establishes correspondences between the feature space of the compressed data and the concept space of the knowledge domain. This mapping is helpful for translating between the statistical patterns identified in the data and the semantic concepts in the knowledge resources. Semantic mapper 1810 may utilize techniques such as zero-shot or few-shot learning to map previously unseen data features to known concepts. It may also employ ontology alignment approaches to establish correspondences between different conceptual frameworks. By creating this semantic bridge, semantic mapper 1810 enables the system to interpret the statistical patterns in the data in terms of meaningful domain concepts, facilitating more informed restoration decisions.

[0146] Rule-based constrainer 1820 applies logical rules and constraints derived from domain knowledge to guide the restoration process. These rules might include but are not limited to physical laws, logical relationships, or domain-specific heuristics that govern valid data configurations. Rule-based constrainer 1820 can identify and filter out statistically plausible but domain-invalid restoration candidates, ensuring that the restored data adheres to known constraints. This component may employ techniques from symbolic AI, such as but not limited to rule engines, logic programming, or constraint satisfaction algorithms, integrated with the neural network architecture through differentiable approximations or hybrid neuro-symbolic approaches.

[0147] Knowledge-feature aligner 1830 performs the task of aligning and combining the information from the data features and the knowledge sources. This alignment process

resolves potential contradictions between statistical evidence and domain knowledge, determining how to optimally weight and integrate different information sources. Knowledge-feature aligner **1830** may employ attention mechanisms to dynamically adjust the influence of different knowledge sources based on their relevance to specific data regions. It may also use uncertainty-aware fusion techniques that consider the reliability of both the statistical patterns and the knowledge sources when integrating them. The output of knowledge-feature aligner **1830** is a unified representation that incorporates both data-driven and knowledge-driven insights.

[0148] Confidence estimator **1840** assesses the reliability and certainty of the integrated knowledge and features. This component quantifies the system's confidence in different aspects of the restoration, taking into account factors such as the consistency between statistical patterns and domain knowledge, the specificity and relevance of the applied knowledge, and the inherent uncertainty in the compressed representation. Confidence estimator **1840** may employ Bayesian techniques, ensemble methods, or dedicated uncertainty estimation networks to generate calibrated confidence scores. These confidence metrics are valuable both for guiding the subsequent restoration process and for providing transparency to end users about the reliability of different restored elements.

[0149] The components of the knowledge integrator **1720** work together in a coordinated manner to achieve effective knowledge integration. Knowledge graph processor **1800** extracts relevant knowledge from the external sources, which semantic mapper **1810** then relates to the data features. Rule-based constrainer **1820** applies domain constraints to ensure validity, while knowledge-feature aligner **1830** combines the statistical and knowledge-based insights.

[0150] Confidence estimator **1840** quantifies the reliability of the integrated representation. This architecture enables several capabilities that enhance the restoration process. The architecture allows the system to incorporate diverse forms of knowledge, from graph-structured conceptual relationships to logical rules and constraints. It provides mechanisms to resolve potential conflicts between statistical evidence and domain knowledge. Additionally, the architecture generates calibrated confidence estimates that can guide both the restoration process and subsequent decision-making. Fourth, it creates a foundation for explainable restoration by tracking how different knowledge sources influence the restoration decisions.

[0151] The modular design of knowledge integrator **1720** also provides flexibility for adaptation to different domains and knowledge representation formats. Components can be modified or replaced to accommodate different types of knowledge sources or to optimize for specific application requirements without necessitating changes to the overall architecture.

[0152] FIG. 19 is a flow diagram illustrating an exemplary method for knowledge-enhanced data compression and restoration using hierarchical autoencoders and correlation networks. In a first step **1900**, domain-specific knowledge sources are collected and integrated into a unified knowledge repository. This initial step involves gathering various forms of domain knowledge, which may include ontologies, knowledge graphs, rule systems, and other structured information sources relevant to the data domain. These diverse knowledge sources are consolidated and harmonized into a

unified repository that can be efficiently accessed and utilized in subsequent steps. This knowledge integration may involve resolving terminological differences, establishing mappings between different knowledge representations, and creating a coherent semantic framework for the domain.

[0153] In a step **1910**, a plurality of data sets is preprocessed and analyzed to identify semantically important features based on the domain-specific knowledge. This preprocessing step goes beyond traditional data cleaning and normalization by incorporating semantic understanding derived from the domain knowledge. The analysis identifies features that are particularly meaningful or significant within the domain context, such as structures, patterns, or relationships that domain experts would recognize as important. This semantic awareness enables more intelligent prioritization during the subsequent compression process.

[0154] In a step **1920**, the analyzed data sets are compressed into compact representations using a semantic-aware hierarchical encoder that prioritizes important domain-specific features. This compression step utilizes a hierarchical autoencoder structure that captures features at multiple scales or levels of abstraction. Unlike conventional compression approaches that might allocate bits based solely on statistical frequency or variance, this semantic-aware encoder dynamically allocates more representational capacity to features identified as semantically important in the previous step. This ensures that domain-significant information is preserved even under high compression ratios.

[0155] In a step **1930**, the compressed data sets are decompressed using a hierarchical decoder to obtain initial decompressed data sets. This hierarchical decoder mirrors the structure of the encoder but operates in reverse, progressively reconstructing the data from its compressed representation. The hierarchical approach allows for the restoration of features at multiple scales, from coarse structures to fine details. While this initial decompression recovers significant portions of the original data, some information loss typically remains due to the compression process.

[0156] In a step **1940**, correlations between the decompressed data sets are identified by leveraging domain knowledge to recognize semantically meaningful relationships. This step enhances the conventional correlation analysis by incorporating semantic understanding. Rather than identifying correlations based solely on statistical patterns, the system leverages domain knowledge to recognize relationships that have semantic significance within the domain context. This might include causal relationships, structural dependencies, or functional associations that are known to exist in the domain.

[0157] In a step **1950**, the decompressed data sets are enhanced using a knowledge graph-enhanced correlation network that applies domain-specific constraints to recover lost information. This restoration step goes beyond statistical inference by explicitly applying domain constraints and relationships represented in the knowledge graph. The knowledge graph-enhanced correlation network can infer missing or degraded information based on known domain patterns, constraints, and relationships. For instance, if certain features are known to always co-occur in valid domain instances, this constraint can guide the restoration of one feature based on the presence of others.

[0158] In a step **1960**, the knowledge-enhanced restored data sets and accompanying explanatory metadata are outputted, indicating which knowledge sources influenced spe-

cific restorations and their confidence levels. This final step not only provides the enhanced restored data but also offers transparency into the restoration process. The explanatory metadata documents which aspects of the domain knowledge were utilized to inform specific restoration decisions and quantifies the confidence or certainty associated with each restored element. This explainability aspect is particularly valuable in applications where understanding the basis for restoration decisions is as important as the restoration itself.

[0159] The method represents a comprehensive approach to data compression and restoration that systematically incorporates domain knowledge at each stage of the process. By leveraging semantic understanding and domain constraints, this method achieves superior restoration quality compared to purely statistical approaches, particularly for complex, structured data where domain knowledge provides valuable context for accurate restoration.

[0160] FIG. 20 is a flow diagram illustrating an alternative exemplary method for semantically guided multi-level data compression and knowledge-based restoration using hierarchical autoencoders and knowledge graph correlation networks. In a first step **2000**, domain knowledge bases, ontologies, rule systems, and semantic dictionaries are integrated into a knowledge graph structure. This initial step involves transforming various forms of structured knowledge into a unified knowledge graph representation that captures concepts as nodes and their relationships as edges. This integration may involve ontology alignment techniques to harmonize different knowledge sources, graph construction from structured rule sets, and incorporation of semantic dictionaries to enhance concept descriptions. The resulting knowledge graph provides a comprehensive and navigable representation of domain knowledge that can be efficiently queried and leveraged in subsequent processing steps.

[0161] In a step **2010**, a plurality of input data sets is normalized and compressed into multi-scale representations using hierarchical encoders that adaptively allocate bits based on semantic importance. This step combines data normalization with semantically guided compression. The hierarchical encoders operate at multiple scales, capturing features ranging from coarse global structures to fine local details. Importantly, the bit allocation during compression is not uniform but is adaptively determined based on the semantic importance of different features as identified through the domain knowledge. This approach ensures that semantically significant features receive more representational capacity, even if they might not be statistically prominent in the data.

[0162] In a step **2020**, the compressed representations are decompressed through hierarchical decoders while preserving semantic relationships between features at different scales. This decompression process utilizes hierarchical decoders that mirror the structure of the encoders but operate in reverse. An aspect of this step is the focus on preserving semantic relationships between features at different scales during decompression. By maintaining these cross-scale semantic relationships, the system can reconstruct more coherent and domain-valid outputs compared to decoders that treat each scale independently.

[0163] In a step **2030**, decompressed outputs are organized and grouped based on semantic relationships identified from the knowledge graph. This organization step goes beyond simple clustering based on statistical similarity. Instead, it

leverages the semantic relationships encoded in the knowledge graph to group decompressed outputs in ways that reflect meaningful domain structures or categories. This semantically informed organization facilitates more effective restoration by providing context for interpreting and enhancing the decompressed data.

[0164] In a step **2040**, relevant domain constraints and relationships applicable to the current data context are extracted from the knowledge graph. Rather than applying all domain knowledge indiscriminately, this step involves selective extraction of constraints and relationships that are specifically relevant to the current data context. This context-aware knowledge extraction improves efficiency and focuses the restoration process on the most pertinent domain insights. The extracted knowledge might include physical laws, logical dependencies, taxonomic relationships, or other domain-specific patterns that can guide the restoration process.

[0165] In a step **2050**, data lost during the compression-decompression process is restored using a knowledge-guided correlation process that infers likely values based on both statistical patterns and domain constraints. This restoration step represents a hybrid approach that combines data-driven and knowledge-driven inference. Statistical patterns within and across the decompressed data sets provide one source of evidence for restoring lost information. These patterns are complemented by domain constraints and relationships extracted from the knowledge graph, which provide additional guidance for plausible restorations. By integrating these different sources of information, the system can restore data more accurately than would be possible using either approach alone.

[0166] In a step **2060**, confidence scores and explanatory metadata linking specific knowledge sources to restoration decisions are generated. This step adds a crucial dimension of transparency to the restoration process. For each restored element, the system generates a confidence score that quantifies the certainty or reliability of the restoration. Additionally, the system creates explanatory metadata that documents which knowledge sources influenced each restoration decision and how they were applied. This information enables users to understand not just what was restored but why and how it was restored, and with what degree of confidence.

[0167] In a step **2070**, the semantically enhanced and restored data sets are output along with the explanatory information. The final output includes both the restored data sets, which benefit from the knowledge-enhanced restoration process, and the accompanying explanatory information generated in the previous step. This combined output provides users with both high-quality restored data and the context needed to interpret and appropriately utilize that data in downstream applications.

[0168] The method represents a comprehensive approach to semantically guided data compression and knowledge-based restoration. By thoroughly integrating domain knowledge at each stage of the process, from initial compression to final output generation, this method achieves superior restoration quality and provides valuable transparency compared to conventional approaches. The emphasis on multi-scale representations, semantic relationships, and explainable restoration makes this method particularly suitable for complex data domains where domain knowledge provides essential context for accurate restoration.

Exemplary Computing Environment

[0169] FIG. 21 illustrates an exemplary computing environment on which an embodiment described herein may be implemented, in full or in part. This exemplary computing environment describes computer-related components and processes supporting enabling disclosure of computer-implemented embodiments. Inclusion in this exemplary computing environment of well-known processes and computer components, if any, is not a suggestion or admission that any embodiment is no more than an aggregation of such processes or components. Rather, implementation of an embodiment using processes and components described in this exemplary computing environment will involve programming or configuration of such processes and components resulting in a machine specially programmed or configured for such implementation. The exemplary computing environment described herein is only one example of such an environment and other configurations of the components and processes are possible, including other relationships between and among components, and/or absence of some processes or components described. Further, the exemplary computing environment described herein is not intended to suggest any limitation as to the scope of use or functionality of any embodiment implemented, in whole or in part, on components or processes described herein.

[0170] The exemplary computing environment described herein comprises a computing device 10 (further comprising a system bus 11, one or more processors 20, a system memory 30, one or more interfaces 40, one or more non-volatile data storage devices 50), external peripherals and accessories 60, external communication devices 70, remote computing devices 80, and cloud-based services 90.

[0171] System bus 11 couples the various system components, coordinating operation of and data transmission between those various system components. System bus 11 represents one or more of any type or combination of types of wired or wireless bus structures including, but not limited to, memory busses or memory controllers, point-to-point connections, switching fabrics, peripheral busses, accelerated graphics ports, and local busses using any of a variety of bus architectures. By way of example, such architectures include, but are not limited to, Industry Standard Architecture (ISA) busses, Micro Channel Architecture (MCA) busses, Enhanced ISA (EISA) busses, Video Electronics Standards Association (VESA) local busses, a Peripheral Component Interconnects (PCI) busses also known as a Mezzanine busses, or any selection of, or combination of, such busses. Depending on the specific physical implementation, one or more of the processors 20, system memory 30 and other components of the computing device 10 can be physically co-located or integrated into a single physical component, such as on a single chip. In such a case, some or all of system bus 11 can be electrical pathways within a single chip structure.

[0172] Computing device may further comprise externally-accessible data input and storage devices 12 such as compact disc read-only memory (CD-ROM) drives, digital versatile discs (DVD), or other optical disc storage for reading and/or writing optical discs 62; magnetic cassettes, magnetic tape, magnetic disk storage, or other magnetic storage devices; or any other medium which can be used to store the desired content and which can be accessed by the computing device 10. Computing device may further comprise externally-accessible data ports or connections 12 such

as serial ports, parallel ports, universal serial bus (USB) ports, and infrared ports and/or transmitter/receivers. Computing device may further comprise hardware for wireless communication with external devices such as IEEE 1394 (“Firewire”) interfaces, IEEE 802.11 wireless interfaces, BLUETOOTH® wireless interfaces, and so forth. Such ports and interfaces may be used to connect any number of external peripherals and accessories 60 such as visual displays, monitors, and touch-sensitive screens 61, USB solid state memory data storage drives (commonly known as “flash drives” or “thumb drives”) 63, printers 64, pointers and manipulators such as mice 65, keyboards 66, and other devices 67 such as joysticks and gaming pads, touchpads, additional displays and monitors, and external hard drives (whether solid state or disc-based), microphones, speakers, cameras, and optical scanners.

[0173] Processors 20 are logic circuitry capable of receiving programming instructions and processing (or executing) those instructions to perform computer operations such as retrieving data, storing data, and performing mathematical calculations. Processors 20 are not limited by the materials from which they are formed or the processing mechanisms employed therein, but are typically comprised of semiconductor materials into which many transistors are formed together into logic gates on a chip (i.e., an integrated circuit or IC). The term processor includes any device capable of receiving and processing instructions including, but not limited to, processors operating on the basis of quantum computing, optical computing, mechanical computing (e.g., using nanotechnology entities to transfer data), and so forth. Depending on configuration, computing device 10 may comprise more than one processor. For example, computing device 10 may comprise one or more central processing units (CPUs) 21, each of which itself has multiple processors or multiple processing cores, each capable of independently or semi-independently processing programming instructions. Further, computing device 10 may comprise one or more specialized processors such as a graphics processing unit (GPU) 22 configured to accelerate processing of computer graphics and images via a large array of specialized processing cores arranged in parallel.

[0174] System memory 30 is processor-accessible data storage in the form of volatile and/or nonvolatile memory. System memory 30 may be either or both of two types: non-volatile memory and volatile memory. Non-volatile memory 30a is not erased when power to the memory is removed, and includes memory types such as read only memory (ROM), electronically-erasable programmable memory (EEPROM), and rewritable solid state memory (commonly known as “flash memory”). Non-volatile memory 30a is typically used for long-term storage of a basic input/output system (BIOS) 31, containing the basic instructions, typically loaded during computer startup, for transfer of information between components within computing device, or a unified extensible firmware interface (UEFI), which is a modern replacement for BIOS that supports larger hard drives, faster boot times, more security features, and provides native support for graphics and mouse cursors. Non-volatile memory 30a may also be used to store firmware comprising a complete operating system 35 and applications 36 for operating computer-controlled devices. The firmware approach is often used for purpose-specific computer-controlled devices such as appliances and Internet-of-Things (IoT) devices where processing power and

data storage space is limited. Volatile memory **30b** is erased when power to the memory is removed and is typically used for short-term storage of data for processing. Volatile memory **30b** includes memory types such as random-access memory (RAM), and is normally the primary operating memory into which the operating system **35**, applications **36**, program modules **37**, and application data **38** are loaded for execution by processors **20**. Volatile memory **30b** is generally faster than non-volatile memory **30a** due to its electrical characteristics and is directly accessible to processors **20** for processing of instructions and data storage and retrieval. Volatile memory **30b** may comprise one or more smaller cache memories which operate at a higher clock speed and are typically placed on the same IC as the processors to improve performance.

[0175] Interfaces **40** may include, but are not limited to, storage media interfaces **41**, network interfaces **42**, display interfaces **43**, and input/output interfaces **44**. Storage media interface **41** provides the necessary hardware interface for loading data from non-volatile data storage devices **50** into system memory **30** and storage data from system memory **30** to non-volatile data storage device **50**. Network interface **42** provides the necessary hardware interface for computing device **10** to communicate with remote computing devices **80** and cloud-based services **90** via one or more external communication devices **70**. Display interface **43** allows for connection of displays **61**, monitors, touchscreens, and other visual input/output devices. Display interface **43** may include a graphics card for processing graphics-intensive calculations and for handling demanding display requirements. Typically, a graphics card includes a graphics processing unit (GPU) and video RAM (VRAM) to accelerate display of graphics. One or more input/output (I/O) interfaces **44** provide the necessary support for communications between computing device **10** and any external peripherals and accessories **60**. For wireless communications, the necessary radio-frequency hardware and firmware may be connected to I/O interface **44** or may be integrated into I/O interface **44**.

[0176] Non-volatile data storage devices **50** are typically used for long-term storage of data. Data on non-volatile data storage devices **50** is not erased when power to the non-volatile data storage devices **50** is removed. Non-volatile data storage devices **50** may be implemented using any technology for non-volatile storage of content including, but not limited to, CD-ROM drives, digital versatile discs (DVD), or other optical disc storage; magnetic cassettes, magnetic tape, magnetic disc storage, or other magnetic storage devices; solid state memory technologies such as EEPROM or flash memory; or other memory technology or any other medium which can be used to store data without requiring power to retain the data after it is written. Non-volatile data storage devices **50** may be non-removable from computing device **10** as in the case of internal hard drives, removable from computing device **10** as in the case of external USB hard drives, or a combination thereof, but computing device will typically comprise one or more internal, non-removable hard drives using either magnetic disc or solid state memory technology. Non-volatile data storage devices **50** may store any type of data including, but not limited to, an operating system **51** for providing low-level and mid-level functionality of computing device **10**, applications **52** for providing high-level functionality of computing device **10**, program modules **53** such as contain-

erized programs or applications, or other modular content or modular programming, application data **54**, and databases **55** such as relational databases, non-relational databases, object oriented databases, BOSQL databases, and graph databases.

[0177] Applications (also known as computer software or software applications) are sets of programming instructions designed to perform specific tasks or provide specific functionality on a computer or other computing devices. Applications are typically written in high-level programming languages such as C++, Java, and Python, which are then either interpreted at runtime or compiled into low-level, binary, processor-executable instructions operable on processors **20**. Applications may be containerized so that they can be run on any computer hardware running any known operating system. Containerization of computer software is a method of packaging and deploying applications along with their operating system dependencies into self-contained, isolated units known as containers. Containers provide a lightweight and consistent runtime environment that allows applications to run reliably across different computing environments, such as development, testing, and production systems.

[0178] The memories and non-volatile data storage devices described herein do not include communication media. Communication media are means of transmission of information such as modulated electromagnetic waves or modulated data signals configured to transmit, not store, information. By way of example, and not limitation, communication media includes wired communications such as sound signals transmitted to a speaker via a speaker wire, and wireless communications such as acoustic waves, radio frequency (RF) transmissions, infrared emissions, and other wireless media.

[0179] External communication devices **70** are devices that facilitate communications between computing device and either remote computing devices **80**, or cloud-based services **90**, or both.

[0180] External communication devices **70** include, but are not limited to, data modems **71** which facilitate data transmission between computing device and the Internet **75** via a common carrier such as a telephone company or internet service provider (ISP), routers **72** which facilitate data transmission between computing device and other devices, and switches **73** which provide direct data communications between devices on a network. Here, modem **71** is shown connecting computing device **10** to both remote computing devices **80** and cloud-based services **90** via the Internet **75**. While modem **71**, router **72**, and switch **73** are shown here as being connected to network interface **42**, many different network configurations using external communication devices **70** are possible. Using external communication devices **70**, networks may be configured as local area networks (LANs) for a single location, building, or campus, wide area networks (WANs) comprising data networks that extend over a larger geographical area, and virtual private networks (VPNs) which can be of any size but connect computers via encrypted communications over public networks such as the Internet **75**. As just one exemplary network configuration, network interface **42** may be connected to switch **73** which is connected to router **72** which is connected to modem **71** which provides access for computing device **10** to the Internet **75**. Further, any combination of wired **77** or wireless **76** communications between and

among computing device 10, external communication devices 70, remote computing devices 80, and cloud-based services 90 may be used. Remote computing devices 80, for example, may communicate with computing device 10 through a variety of communication channels 74 such as through switch 73 via a wired 77 connection, through router 72 via a wireless connection 76, or through modem 71 via the Internet 75. Furthermore, while not shown here, other hardware that is specifically designed for servers may be employed. For example, secure socket layer (SSL) acceleration cards can be used to offload SSL encryption computations, and transmission control protocol/internet protocol (TCP/IP) offload hardware and/or packet classifiers on network interfaces 42 may be installed and used at server devices.

[0181] In a networked environment, certain components of computing device 10 may be fully or partially implemented on remote computing devices 80 or cloud-based services 90. Data stored in non-volatile data storage device 50 may be received from, shared with, duplicated on, or offloaded to a non-volatile data storage device on one or more remote computing devices 80 or in a cloud computing service 92. Processing by processors 20 may be received from, shared with, duplicated on, or offloaded to processors of one or more remote computing devices 80 or in a distributed computing service 93. By way of example, data may reside on a cloud computing service 92, but may be usable or otherwise accessible for use by computing device 10. Also, certain processing subtasks may be sent to a microservice 91 for processing with the result being transmitted to computing device 10 for incorporation into a larger processing task. Also, while components and processes of the exemplary computing environment are illustrated herein as discrete units (e.g., OS 51 being stored on non-volatile data storage device 51 and loaded into system memory 35 for use) such processes and components may reside or be processed at various times in different components of computing device 10, remote computing devices 80, and/or cloud-based services 90.

[0182] In an implementation, the disclosed systems and methods may utilize, at least in part, containerization techniques to execute one or more processes and/or steps disclosed herein. Containerization is a lightweight and efficient virtualization technique that allows you to package and run applications and their dependencies in isolated environments called containers. One of the most popular containerization platforms is Docker, which is widely used in software development and deployment. Containerization, particularly with open-source technologies like Docker and container orchestration systems like Kubernetes, is a common approach for deploying and managing applications. Containers are created from images, which are lightweight, standalone, and executable packages that include application code, libraries, dependencies, and runtime. Images are often built from a Dockerfile or similar, which contains instructions for assembling the image. Dockerfiles are configuration files that specify how to build a Docker image. Systems like Kubernetes also support containerd or CRI-O. They include commands for installing dependencies, copying files, setting environment variables, and defining runtime configurations. Docker images are stored in repositories, which can be public or private. Docker Hub is an exemplary public registry, and organizations often set up private registries for security and version control using tools such as

Hub, JFrog Artifactory and Bintray, Github Packages or Container registries. Containers can communicate with each other and the external world through networking. Docker provides a bridge network by default, but can be used with custom networks. Containers within the same network can communicate using container names or IP addresses.

[0183] Remote computing devices 80 are any computing devices not part of computing device 10. Remote computing devices 80 include, but are not limited to, personal computers, server computers, thin clients, thick clients, personal digital assistants (PDAs), mobile telephones, watches, tablet computers, laptop computers, multiprocessor systems, microprocessor based systems, set-top boxes, programmable consumer electronics, video game machines, game consoles, portable or handheld gaming units, network terminals, desktop personal computers (PCs), minicomputers, main frame computers, network nodes, virtual reality or augmented reality devices and wearables, and distributed or multiprocessor computing environments. While remote computing devices 80 are shown for clarity as being separate from cloud-based services 90, cloud-based services 90 are implemented on collections of networked remote computing devices 80.

[0184] Cloud-based services 90 are Internet-accessible services implemented on collections of networked remote computing devices 80. Cloud-based services are typically accessed via application programming interfaces (APIs) which are software interfaces which provide access to computing services within the cloud-based service via API calls, which are pre-defined protocols for requesting a computing service and receiving the results of that computing service. While cloud-based services may comprise any type of computer processing or storage, three common categories of cloud-based services 90 are microservices 91, cloud computing services 92, and distributed computing services 93.

[0185] Microservices 91 are collections of small, loosely coupled, and independently deployable computing services. Each microservice represents a specific computing functionality and runs as a separate process or container. Microservices promote the decomposition of complex applications into smaller, manageable services that can be developed, deployed, and scaled independently. These services communicate with each other through well-defined application programming interfaces (APIs), typically using lightweight protocols like HTTP, gRPC, or message queues such as Kafka. Microservices 91 can be combined to perform more complex processing tasks.

[0186] Cloud computing services 92 are delivery of computing resources and services over the Internet 75 from a remote location. Cloud computing services 92 provide additional computer hardware and storage on as-needed or subscription basis. Cloud computing services 92 can provide large amounts of scalable data storage, access to sophisticated software and powerful server-based processing, or entire computing infrastructures and platforms. For example, cloud computing services can provide virtualized computing resources such as virtual machines, storage, and networks, platforms for developing, running, and managing applications without the complexity of infrastructure management, and complete software applications over the Internet on a subscription basis.

[0187] Distributed computing services 93 provide large-scale processing using multiple interconnected computers or

nodes to solve computational problems or perform tasks collectively. In distributed computing, the processing and storage capabilities of multiple machines are leveraged to work together as a unified system. Distributed computing services are designed to address problems that cannot be efficiently solved by a single computer or that require large-scale computational power. These services enable parallel processing, fault tolerance, and scalability by distributing tasks across multiple nodes.

[0188] Although described above as a physical device, computing device **10** can be a virtual computing device, in which case the functionality of the physical components herein described, such as processors **20**, system memory **30**, network interfaces **40**, and other like components can be provided by computer-executable instructions. Such computer-executable instructions can execute on a single physical computing device, or can be distributed across multiple physical computing devices, including being distributed across multiple physical computing devices in a dynamic manner such that the specific, physical computing devices hosting such computer-executable instructions can dynamically change over time depending upon need and availability. In the situation where computing device **10** is a virtualized device, the underlying physical computing devices hosting such a virtualized computing device can, themselves, comprise physical components analogous to those described above, and operating in a like manner. Furthermore, virtual computing devices can be utilized in multiple layers with one virtual computing device executing within the construct of another virtual computing device. Thus, computing device **10** may be either a physical computing device or a virtualized computing device within which computer-executable instructions can be executed in a manner consistent with their execution by a physical computing device. Similarly, terms referring to physical components of the computing device, as utilized herein, mean either those physical components or virtualizations thereof performing the same or equivalent functions. The skilled person will be aware of a range of possible modifications of the various aspects described above. Accordingly, the present invention is defined by the claims and their equivalents.

What is claimed is:

1. A computer system comprising:

a hardware memory, wherein the computer system is configured to execute software instructions stored on nontransitory machine-readable storage media that:

train a correlation network using sets of cross-correlated training data sets;

obtain a plurality of input data sets;

compress the input data sets into compressed data sets using a plurality of encoders within a plurality of hierarchical autoencoders;

decompress the compressed data sets using a plurality of decoders within the plurality of hierarchical autoencoders to obtain decompressed data sets; and

enhance the decompressed data sets using a restoration process that incorporates external information sources to recover lost or degraded information, thereby generating enhanced restored data sets.

2. The system of claim **1**, wherein the data sets comprise data organized by type prior to preprocessing.

3. The system of claim **1**, wherein the data sets comprise spectral data.

4. The computer system of claim **1**, wherein the restoration process utilizes a correlation network that leverages semantic relationships within the data to enhance restoration quality.

5. A method for compressing and restoring data, comprising the steps of:

training a correlation network using sets of cross-correlated training data sets;

obtaining a plurality of input data sets;

compressing the input data sets into compressed data sets using a plurality of encoders within a plurality of hierarchical autoencoders;

decompressing the compressed data sets using a plurality of decoders within the plurality of hierarchical autoencoders to obtain decompressed data sets; and

enhancing the decompressed data sets using a restoration process that incorporates external information sources to recover lost or degraded information, thereby generating enhanced restored data sets.

6. The method of claim **5**, wherein the data sets comprise data organized by type prior to preprocessing.

7. The method of claim **5**, wherein the data sets comprise spectral data.

8. The method of claim **5**, wherein the restoration process utilizes a correlation network that leverages semantic relationships within the data to enhance restoration quality.

* * * * *