

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent Application Publication

20250260584

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

Bulusu; Vijay et al.

---

### Cloud-Based Man-in-the-Middle Inspection of Encrypted Traffic

---

#### Abstract

A method for inspecting encrypted network traffic in a cloud-based security system is provided. A node receives a request from a user device targeting a server and obtains a domain certificate corresponding to the server. The method establishes a first encrypted tunnel between the user device and the node, and a second encrypted tunnel between the node and the server using the obtained certificate. The encrypted traffic flowing between the user device and the server is inspected at the node. The method leverages a cloud-based hardware security module (HSM) to securely generate and store intermediate certificate authority keys compliant with FIPS 140-2 Level 3 standards, facilitating secure man-in-the-middle (MITM) inspection. The method also enables caching and synchronization of domain certificates across distributed nodes, providing scalable and secure traffic monitoring.

---

**Inventors:** Bulusu; Vijay (Fremont, CA), Maheshwari; Akshat (Bangalore, IN), Singh; Harpreet (San Jose, CA), Kumar; Sujay (Bangalore, IN), Pergament; Lidor (San Mateo, CA), Devarajan; Srikanth (Cupertino, CA)

**Applicant:** Zscaler, Inc. (San Jose, CA)

**Family ID:** 83509663

**Assignee:** Zscaler, Inc. (San Jose, CA)

**Appl. No.:** 19/174342

**Filed:** April 09, 2025

#### Related U.S. Application Data

parent US continuation 17843095 20220617 parent-grant-document US 12309295 child US 19174342

parent US continuation 16863475 20200430 parent-grant-document US 11394563 child US 17715137

## Publication Classification

**Int. Cl.:** H04L9/32 (20060101); H04L9/40 (20220101)

**U.S. Cl.:**

**CPC** H04L9/3263 (20130101); H04L63/0428 (20130101); H04L63/1416 (20130101);

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION(S) [0001] The present disclosure is a continuation of U.S. patent application Ser. No. 17/843,095, filed Jun. 17, 2022, and entitled “Cloud-based man-in-the-middle inspection of encrypted traffic using cloud-based multi-tenant HSM infrastructure,” which is a continuation-in-part of U.S. patent application Ser. No. 17/715,137, filed Apr. 7, 2022, now U.S. Pat. No. 11,750,405, issued Sep. 5, 2023, and entitled “Encrypted traffic inspection in a cloud-based security system,” which is a continuation of U.S. patent application Ser. No. 16/863,475, filed Apr. 30, 2020, now U.S. Pat. No. 11,394,563, issued Jul. 19, 2022, and entitled “Encrypted traffic inspection in a cloud-based security system,” the contents of each are incorporated by reference in their entirety.

### FIELD OF THE DISCLOSURE

[0002] The present disclosure relates generally to networking and computing. More particularly, the present disclosure relates to systems and methods for cloud-based hardware security modules (HSM) with multi-tenant key infrastructure for man-in-the-middle inspection of encrypted traffic.

### BACKGROUND OF THE DISCLOSURE

[0003] There is a significant amount of encrypted traffic on the Internet. For example, protocols such as SSL, TLS, DTLS, HTTPS, etc. are used to provide privacy and data integrity. According to some forecasts, 70% or more of all Web traffic now uses SSL, and these numbers are growing. Encrypted traffic presents a security hole, i.e., a blind spot. Enterprises conventionally have deployed appliances and other devices at the network perimeter to perform security functions. In terms of encrypted traffic, the appliances need to break the encryption in order to monitor the traffic. This is resource intense, and conventional appliances simply do not scale. As such, most enterprises simply forego the inspection of encrypted traffic. Other studies have shown that the majority of malware today is hidden in encrypted traffic. Also, encrypted traffic presents a problem in terms of Data Loss Prevention (DLP) because sensitive data is typically concealed in SSL/TLS traffic, which is difficult and expensive to inspect (in terms of cost, processing capability, and latency). Without visibility and control, organizations are at an increased risk of data loss, due either to unintentional or malicious reasons. The conventional appliance and network perimeter security approach is breaking down with the mobility of users, the processing capability of user devices, etc. As such, security is moving to the cloud, namely as a service offered through a cloud-based system.

[0004] A hardware security module (HSM) is a physical computing device that safeguards and manages digital keys, performs encryption and decryption functions for digital signatures, strong authentication and other cryptographic functions. These modules traditionally come in the form of a plug-in card or an external device that attaches directly to a computer or network server. A hardware security module contains one or more secure cryptoprocessor chips There are various use

cases for an HSM such as a central authority (CA), payments, secure tunnel establishment, cryptocurrency, etc.

[0005] In an example application, most large and major enterprises prefer to bring their own intermediate CA to a cloud-based security system in order to perform man-in-the-middle (MITM) based transport layer security (TLS) inspection. In this model, the cloud-based security system essentially acts as an external, to the enterprise “boundary,” intermediate CA of the enterprise public key infrastructure (PKI) and generates/stores an intermediate CA private key in the cloud boundary on disk (e.g., structured query language (SQL) database (DB)). Disadvantageously, storing private keys for a customer on a SQL DB without hardware protection leaves them prone to be exploited and exposed. Whether the exposure is driven by a malicious intruder, malicious insider/operator or even unintentional, the consequence is that an adversary who controls the key can sign TLS certificates for malicious destinations (e.g., credential stealing sites) on behalf of the customer, which would be blindly trusted by every user that is configured to their root CA. For large organizations and financial service providers in particular, this risk is unacceptable and makes them out of compliant with PKI governance policies that mandate that any certificate private key issued from an Enterprise PKI (Root or iCA) that has the power to issue subordinate certificates of any kind must be generated in and persist its entire lifetime in a FIPS 140-2 Level 3 HSM.

#### BRIEF SUMMARY OF THE DISCLOSURE

[0006] The present disclosure relates to systems and methods for man-in-the-middle inspection of encrypted traffic using cloud-based hardware security modules (HSM) with multi-tenant key infrastructure. Of note, any certificate private key issued from an Enterprise PKI (Root or iCA) that has the power to issue subordinate certificates of any kind must be generated in and persist for its entire lifetime in a FIPS 140-2 Level 3 HSM. The Federal Information Processing Standard Publication 140-2, (FIPS PUB 140-2), the contents of which are incorporated by reference herein, is a U.S. government computer security standard used to approve cryptographic modules. The title is Security Requirements for Cryptographic Modules. Initial publication was on May 25, 2001, and was last updated Dec. 3, 2002.

[0007] Also, the present disclosure relates to systems and methods for encrypted traffic inspection in a cloud-based security system, such as Secure Sockets Layer (SSL), Transport Layer Security (TLS), Hypertext Transfer Protocol Secure (HTTPS), and the like. The cloud-based security system is configured to monitor users in an inline manner as a proxy or Secure Web or Internet Gateway, including monitoring encrypted traffic, e.g., Secure Sockets Layer (SSL)/Transport Layer Security (TLS) traffic. Based on this proxy by the design aspect of the cloud-based security system, the cloud-based security system can provide inspection on encrypted traffic, such as SSL, TLS, DTLS, HTTPS, etc., without the inspection limitations of appliances. Various approaches are contemplated, including a snooping approach, a Man-in-the-Middle (MitM) proxy approach, and the like. The snooping approach includes snooping session keys and utilizing the snooped keys to non-intrusively monitor the encrypted traffic. Advantageously, this approach does not terminate the encrypted traffic. The MitM proxy approach has a cloud node that sits as a proxy between a user device and an endpoint where the proxy breaks the encrypted traffic in the middle. With the inspection of encrypted traffic, the cloud-based security system can perform a full suite of security functions on the traffic.

---

## Description

#### BRIEF DESCRIPTION OF THE DRAWINGS

[0008] The present disclosure is illustrated and described herein with reference to the various drawings, in which like reference numbers are used to denote like system components/method steps, as appropriate, and in which:

[0009] FIG. 1 is a network diagram of a cloud-based system offering security as a service.

[0010] FIG. 2 is a network diagram of an example implementation of the cloud-based system.

[0011] FIG. 3 is a block diagram of a server that may be used in the cloud-based system of FIGS. 1 and 2 or the like.

[0012] FIG. 4 is a block diagram of a user device that may be used with the cloud-based system of FIGS. 1 and 2 or the like.

[0013] FIG. 5 is a network diagram of the cloud-based system illustrating an application on user devices with users configured to operate through the cloud-based system.

[0014] FIG. 6 is a network diagram of a Zero Trust Network Access (ZTNA) application utilizing the cloud-based system of FIGS. 1 and 2.

[0015] FIG. 7 is a network diagram of the cloud-based system of FIGS. 1 and 2 in an application of digital experience monitoring.

[0016] FIG. 8 is a flow diagram illustrating an example handshake for HTTPS to describe a secure, encrypted tunnel between a client (e.g., the user device of FIG. 4) and a server (e.g., the server of FIG. 3).

[0017] FIG. 9 is a screenshot of packet capture showing SSL packets as they are exchanged between a client and a server.

[0018] FIG. 10 is a flow diagram illustrating an embodiment of SSL inspection with the cloud-based system as a proxy.

[0019] FIG. 11 is a flow diagram of details of an SSL handshake process between an SSL client and an SSL server.

[0020] FIG. 12 is a flow diagram of a process performing SSL interception through an interception proxy in the handshake process.

[0021] FIG. 13 is a network diagram of a network with an enforcement node operating as an interception proxy.

[0022] FIG. 14 is a network diagram of a network with the enforcement node operating as a snooping proxy to perform SSL interception without breaking the tunnel as with the interception proxy.

[0023] FIG. 15 is a flowchart of a process for SSL (or other type of encrypted traffic) inspection by snooping, such as via a node operating as the snooping proxy.

[0024] FIG. 16 is a diagram of encrypted tunnel inspection with the cloud-based system.

[0025] FIG. 17 is a diagram of key storage and key distribution, in an embodiment, in the cloud-based system.

[0026] FIG. 18 is a diagram illustrating certificates, their descriptions, and private key storage, based on FIG. 17.

[0027] FIG. 19 is a network diagram of the cloud-based system and a cloud HSM system.

[0028] FIG. 20 is a diagram illustrating workflow between a user device, an enforcement node, and a cloud HSM system.

[0029] FIG. 21 is a diagram of interactions between the cloud-based system and the cloud HSM.

[0030] FIG. 22 is a diagram illustrating certificates, their descriptions, and private key storage, based on FIG. 21.

[0031] FIG. 23 is a diagram of key storage and key distribution with the cloud HSM, in an embodiment, in the cloud-based system.

[0032] FIG. 24 is a flowchart of a process for man-in-the-middle inspection of encrypted traffic with cloud-based hardware security modules (HSM) with multi-tenant key infrastructure.

#### DETAILED DESCRIPTION OF THE DISCLOSURE

[0033] Again, the present disclosure relates to systems and methods for man-in-the-middle inspection of encrypted traffic with cloud-based hardware security modules (HSM) with multi-tenant key infrastructure. Of note, any certificate private key issued from an Enterprise PKI (Root or iCA) that has the power to issue subordinate certificates of any kind must be generated in and

persist for its entire lifetime in a FIPS 140-2 Level 3 HSM. The Federal Information Processing Standard Publication 140-2, (FIPS PUB 140-2), the contents of which are incorporated by reference herein, is a U.S. government computer security standard used to approve cryptographic modules. The title is Security Requirements for Cryptographic Modules. Initial publication was on May 25, 2001, and was last updated Dec. 3, 2002.

#### § 1.0 Example Cloud-Based System Architecture

[0034] FIG. 1 is a network diagram of a cloud-based system **100** offering security as a service. Specifically, the cloud-based system **100** can offer a Secure Internet and Web Gateway as a service to various users **102**, as well as other cloud services. In this manner, the cloud-based system **100** is located between the users **102** and the Internet as well as any cloud services **106** (or applications) accessed by the users **102**. As such, the cloud-based system **100** provides inline monitoring inspecting traffic between the users **102**, the Internet **104**, and the cloud services **106**, including secure sockets layer (SSL) traffic. The cloud-based system **100** can offer access control, threat prevention, data protection, etc. The access control can include a cloud-based firewall, cloud-based intrusion detection, uniform resource locator (URL) filtering, bandwidth control, domain name system (DNS) filtering, etc. Threat prevention can include cloud-based intrusion prevention, protection against advanced threats (malware, spam, cross-site scripting (XSS), phishing, etc.), cloud-based sandbox, antivirus, DNS security, etc. The data protection can include data loss prevention (DLP), cloud application security such as via a cloud access security broker (CASB), file type control, etc.

[0035] The cloud-based firewall can provide deep packet inspection (DPI) and access controls across various ports and protocols as well as being application and user aware. The URL filtering can block, allow, or limit website access based on policy for a user, group of users, or entire organization, including specific destinations or categories of URLs (e.g., gambling, social media, etc.). The bandwidth control can enforce bandwidth policies and prioritize critical applications such as relative to recreational traffic. DNS filtering can control and block DNS requests against known and malicious destinations.

[0036] The cloud-based intrusion prevention and advanced threat protection can deliver full threat protection against malicious content such as browser exploits, scripts, identified botnets and malware callbacks, etc. The cloud-based sandbox can block zero-day exploits (just identified) by analyzing unknown files for malicious behavior. Advantageously, the cloud-based system **100** is multi-tenant and can service a large volume of the users **102**. As such, newly discovered threats can be promulgated throughout the cloud-based system **100** for all tenants practically instantaneously. The antivirus protection can include antivirus, antispymware, antimalware, etc. protection for the users **102**, using signatures sourced and constantly updated. The DNS security can identify and route command-and-control connections to threat detection engines for full content inspection.

[0037] The DLP can use standard and/or custom dictionaries to continuously monitor the users **102**, including compressed and/or SSL-encrypted traffic. Again, being in a cloud implementation, the cloud-based system **100** can scale this monitoring with near-zero latency on the users **102**. The cloud application security can include CASB functionality to discover and control user access to known and unknown cloud services **106**. The file type controls enable true file type control by the user, location, destination, etc. to determine which files are allowed or not.

[0038] The cloud-based system **100** can provide other security functions, including, for example, micro-segmentation, workload segmentation, API security, cloud security posture management (CSPM), user identity management, and the like. That is, the cloud-based system **100** provides a network architecture that enables delivery of any cloud-based security service, including emerging frameworks.

[0039] For illustration purposes, the users **102** of the cloud-based system **100** can include a mobile device **110**, a headquarters (HQ) **112** which can include or connect to a data center (DC) **114**, Internet of Things (IoT) devices **116**, a branch office/remote location **118**, etc., and each includes

one or more user devices (an example user device **300** (user equipment (UE)) is illustrated in FIG. 5). The devices **110**, **116**, and the locations **112**, **114**, **118** are shown for illustrative purposes, and those skilled in the art will recognize there are various access scenarios and other users **102** for the cloud-based system **100**, all of which are contemplated herein. The users **102** can be associated with a tenant, which may include an enterprise, a corporation, an organization, etc. That is, a tenant is a group of users who share a common access with specific privileges to the cloud-based system **100**, a cloud service, etc. In an embodiment, the headquarters **112** can include an enterprise's network with resources in the data center **114**. The mobile device **110** can be a so-called road warrior, i.e., users that are off-site, on-the-road, etc. Those skilled in the art will recognize a user **102** has to use a corresponding user device **300** for accessing the cloud-based system **100** and the like, and the description herein may use the user **102** and/or the user device **300** interchangeably. [0040] Further, the cloud-based system **100** can be multi-tenant, with each tenant having its own users **102** and configuration, policy, rules, etc. One advantage of the multi-tenancy and a large volume of users is the zero-day/zero-hour protection in that a new vulnerability can be detected and then instantly remediated across the entire cloud-based system **100**. The same applies to policy, rule, configuration, etc. changes—they are instantly remediated across the entire cloud-based system **100**. As well, new features in the cloud-based system **100** can also be rolled up simultaneously across the user base, as opposed to selective and time-consuming upgrades on every device at the locations **112**, **114**, **118**, and the devices **110**, **116**.

[0041] Logically, the cloud-based system **100** can be viewed as an overlay network between users (at the locations **112**, **114**, **118**, and the devices **110**, **116**) and the Internet **104** and the cloud services **106**. Previously, the IT deployment model included enterprise resources and applications stored within the data center **114** (i.e., physical devices) behind a firewall (perimeter), accessible by employees, partners, contractors, etc. on-site or remote via virtual private networks (VPNs), etc. The cloud-based system **100** is replacing the conventional deployment model. The cloud-based system **100** can be used to implement these services in the cloud without requiring the physical devices and management thereof by enterprise IT administrators. As an ever-present overlay network, the cloud-based system **100** can provide the same functions as the physical devices and/or appliances regardless of geography or location of the users **102**, as well as independent of platform, operating system, network access technique, network access provider, etc.

[0042] There are various techniques to forward traffic between the users **102** at the locations **112**, **114**, **118**, and via the devices **110**, **116**, and the cloud-based system **100**. Typically, the locations **112**, **114**, **118** can use tunneling where all traffic is forward through the cloud-based system **100**. For example, various tunneling protocols are contemplated, such as GRE, L2TP, IPsec, customized tunneling protocols, etc. The devices **110**, **116**, when not at one of the locations **112**, **114**, **118** can use a local application that forwards traffic, a proxy such as via a proxy auto-config (PAC) file, and the like. An application of the local application is the application **350** described in detail herein as a connector application. A key aspect of the cloud-based system **100** is all traffic between the users **102** and the Internet **104** or the cloud services **106** is via the cloud-based system **100**. As such, the cloud-based system **100** has visibility to enable various functions, all of which are performed off the user device in the cloud.

[0043] The cloud-based system **100** can also include a management system **120** for tenant access to provide global policy and configuration as well as real-time analytics. This enables IT administrators to have a unified view of user activity, threat intelligence, application usage, etc. For example, IT administrators can drill-down to a per-user level to understand events and correlate threats, to identify compromised devices, to have application visibility, and the like. The cloud-based system **100** can further include connectivity to an identity provider (IDP) **122** for authentication of the users **102** and to a security information and event management (SIEM) system **124** for event logging. The system **124** can provide alert and activity logs on a per-user **102** basis.

[0044] FIG. 2 is a network diagram of an example implementation of the cloud-based system **100**.

In an embodiment, the cloud-based system **100** includes a plurality of enforcement nodes (EN) **150**, labeled as enforcement nodes **150-1**, **150-2**, **150-N**, interconnected to one another and interconnected to a central authority (CA) **152**. Note, the nodes **150** are called “enforcement” nodes **150** but they can be simply referred to as nodes **150** in the cloud-based system **100**. Also, the nodes **150** can be referred to as service edges. The nodes **150** and the central authority **152**, while described as nodes, can include one or more servers, including physical servers, virtual machines (VM) executed on physical hardware, etc. An example of a server is illustrated in FIG. 4. The cloud-based system **100** further includes a log router **154** that connects to a storage cluster **156** for supporting log maintenance from the enforcement nodes **150**. The central authority **152** provide centralized policy, real-time threat updates, etc. and coordinates the distribution of this data between the enforcement nodes **150**. The enforcement nodes **150** provide an onramp to the users **102** and are configured to execute policy, based on the central authority **152**, for each user **102**. The enforcement nodes **150** can be geographically distributed, and the policy for each user **102** follows that user **102** as he or she connects to the nearest (or other criteria) enforcement node **150**. Of note, the cloud-based system is an external system meaning it is separate from the tenant's private networks (enterprise networks) as well as from networks associated with the devices **110**, **116**, and locations **112**, **118**.

[0045] The enforcement nodes **150** are full-featured secure internet gateways that provide integrated internet security. They inspect all web traffic bi-directionally for malware and enforce security, compliance, and firewall policies, as described herein, as well as various additional functionality. In an embodiment, each enforcement node **150** has two main modules for inspecting traffic and applying policies: a web module and a firewall module. The enforcement nodes **150** are deployed around the world and can handle hundreds of thousands of concurrent users with millions of concurrent sessions. Because of this, regardless of where the users **102** are, they can access the Internet **104** from any device, and the enforcement nodes **150** protect the traffic and apply corporate policies. The enforcement nodes **150** can implement various inspection engines therein, and optionally, send sandboxing to another system. The enforcement nodes **150** include significant fault tolerance capabilities, such as deployment in active-active mode to ensure availability and redundancy as well as continuous monitoring.

[0046] In an embodiment, customer traffic is not passed to any other component within the cloud-based system **100**, and the enforcement nodes **150** can be configured never to store any data to disk. Packet data is held in memory for inspection and then, based on policy, is either forwarded or dropped. Log data generated for every transaction is compressed, tokenized, and exported over secure transport layer security (TLS) connections to the log routers **154** that direct the logs to the storage cluster **156**, hosted in the appropriate geographical region, for each organization. In an embodiment, all data destined for or received from the Internet is processed through one of the enforcement nodes **150**. In another embodiment, specific data specified by each tenant, e.g., only email, only executable files, etc., is processed through one of the enforcement nodes **150**.

[0047] Each of the enforcement nodes **150** may generate a decision vector  $D=[d_1, d_2, \dots, d_n]$  for a content item of one or more parts  $C=[c_1, c_2, \dots, c_m]$ . Each decision vector may identify a threat classification, e.g., clean, spyware, malware, undesirable content, innocuous, spam email, unknown, etc. For example, the output of each element of the decision vector  $D$  may be based on the output of one or more data inspection engines. In an embodiment, the threat classification may be reduced to a subset of categories, e.g., violating, non-violating, neutral, unknown. Based on the subset classification, the enforcement node **150** may allow the distribution of the content item, preclude distribution of the content item, allow distribution of the content item after a cleaning process, or perform threat detection on the content item. In an embodiment, the actions taken by one of the enforcement nodes **150** may be determinative on the threat classification of the content item and on a security policy of the tenant to which the content item is being sent from or from which the content item is being requested by. A content item is violating if, for any part  $C=[c_1, c_2, \dots$

... , cm] of the content item, at any of the enforcement nodes **150**, any one of the data inspection engines generates an output that results in a classification of “violating.”

[0048] The central authority **152** hosts all customer (tenant) policy and configuration settings. It monitors the cloud and provides a central location for software and database updates and threat intelligence. Given the multi-tenant architecture, the central authority **152** is redundant and backed up in multiple different data centers. The enforcement nodes **150** establish persistent connections to the central authority **152** to download all policy configurations. When a new user connects to an enforcement node **150**, a policy request is sent to the central authority **152** through this connection. The central authority **152** then calculates the policies that apply to that user **102** and sends the policy to the enforcement node **150** as a highly compressed bitmap.

[0049] The policy can be tenant-specific and can include access privileges for users, websites and/or content that is disallowed, restricted domains, DLP dictionaries, etc. Once downloaded, a tenant's policy is cached until a policy change is made in the management system **120**. The policy can be tenant-specific and can include access privileges for users, websites and/or content that is disallowed, restricted domains, DLP dictionaries, etc. When this happens, all of the cached policies are purged, and the enforcement nodes **150** request the new policy when the user **102** next makes a request. In an embodiment, the enforcement nodes **150** exchange “heartbeats” periodically, so all enforcement nodes **150** are informed when there is a policy change. Any enforcement node **150** can then pull the change in policy when it sees a new request.

[0050] The cloud-based system **100** can be a private cloud, a public cloud, a combination of a private cloud and a public cloud (hybrid cloud), or the like. Cloud computing systems and methods abstract away physical servers, storage, networking, etc., and instead offer these as on-demand and elastic resources. The National Institute of Standards and Technology (NIST) provides a concise and specific definition which states cloud computing is a model for enabling convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. Cloud computing differs from the classic client-server model by providing applications from a server that are executed and managed by a client's web browser or the like, with no installed client version of an application required. Centralization gives cloud service providers complete control over the versions of the browser-based and other applications provided to clients, which removes the need for version upgrades or license management on individual client computing devices. The phrase “software as a service” (SaaS) is sometimes used to describe application programs offered through cloud computing. A common shorthand for a provided cloud computing service (or even an aggregation of all existing cloud services) is “the cloud.” The cloud-based system **100** is illustrated herein as an example embodiment of a cloud-based system, and other implementations are also contemplated.

[0051] As described herein, the terms cloud services and cloud applications may be used interchangeably. The cloud service **106** is any service made available to users on-demand via the Internet, as opposed to being provided from a company's on-premises servers. A cloud application, or cloud app, is a software program where cloud-based and local components work together. The cloud-based system **100** can be utilized to provide example cloud services, including Zscaler Internet Access (ZIA), Zscaler Private Access (ZPA), and Zscaler Digital Experience (ZDX), all from Zscaler, Inc. (the assignee and applicant of the present application). Also, there can be multiple different cloud-based systems **100**, including ones with different architectures and multiple cloud services. The ZIA service can provide the access control, threat prevention, and data protection described above with reference to the cloud-based system **100**. ZPA can include access control, microservice segmentation, etc. The ZDX service can provide monitoring of user experience, e.g., Quality of Experience (QoE), Quality of Service (QoS), etc., in a manner that can gain insights based on continuous, inline monitoring. For example, the ZIA service can provide a user with Internet Access, and the ZPA service can provide a user with access to enterprise



resources instead of traditional virtual private networks (VPNs), namely ZPA provides zero trust network access (ZTNA). Those of ordinary skill in the art will recognize various other types of cloud services **106** are also contemplated. Also, other types of cloud architectures are also contemplated, with the cloud-based system **100** presented for illustration purposes.

### § 1.1 Private Nodes Hosted by Tenants or Service Providers

[0052] The nodes **150** that service multi-tenant users **102** may be located in data centers. These nodes **150** can be referred to as public nodes **150** or public service edges. In embodiment, the nodes **150** can be located on-premises with tenants (enterprise) as well as service providers. These nodes can be referred to as private nodes **150** or private service edges. In operation, these private nodes **150** can perform the same functions as the public nodes **150**, can communicate with the central authority **152**, and the like. In fact, the private nodes **150** can be considered in the same cloud-based system **100** as the public nodes **150**, except located on-premises. When a private node **150** is located in an enterprise network, the private node **150** can have a single tenant corresponding to the enterprise; of course, the cloud-based system **100** is still multi-tenant, but these particular nodes are serving only a single tenant. When a private node **150** is located in a service provider's network, the private node **150** can be multi-tenant for customers of the service provider. Those skilled in the art will recognize various architectural approaches are contemplated. The cloud-based system **100** is a logical construct providing a security service.

### § 2.0 User Device Application for Traffic Forwarding and Monitoring

[0053] FIG. **3** is a network diagram of the cloud-based system **100** illustrating an application **350** on user devices **300** with users **102** configured to operate through the cloud-based system **100**. Different types of user devices **300** are proliferating, including bring your own device (BYOD) as well as IT-managed devices. The conventional approach for a user device **300** to operate with the cloud-based system **100** as well as for accessing enterprise resources includes complex policies, VPNs, poor user experience, etc. The application **350** can automatically forward user traffic with the cloud-based system **100** as well as ensuring that security and access policies are enforced, regardless of device, location, operating system, or application. The application **350** automatically determines if a user **102** is looking to access the open Internet **104**, a SaaS app, or an internal app running in public, private, or the datacenter and routes mobile traffic through the cloud-based system **100**. The application **350** can support various cloud services, including ZIA, ZPA, ZDX, etc., allowing the best-in-class security with zero trust access to internal apps. As described herein, the application **350** can also be referred to as a connector application.

[0054] The application **350** is configured to auto-route traffic for seamless user experience. This can be protocol as well as application-specific, and the application **350** can route traffic with a nearest or best fit enforcement node **150**. Further, the application **350** can detect trusted networks, allowed applications, etc. and support secure network access. The application **350** can also support the enrollment of the user device **300** prior to accessing applications. The application **350** can uniquely detect the users **102** based on fingerprinting the user device **300**, using criteria like device model, platform, operating system, etc. The application **350** can support mobile device management (MDM) functions, allowing IT personnel to deploy and manage the user devices **300** seamlessly. This can also include the automatic installation of client and SSL certificates during enrollment. Finally, the application **350** provides visibility into device and app usage of the user **102** of the user device **300**.

[0055] The application **350** supports a secure, lightweight tunnel between the user device **300** and the cloud-based system **100**. For example, the lightweight tunnel can be HTTP-based. With the application **350**, there is no requirement for PAC files, an IPsec VPN, authentication cookies, or user **102** setup.

### § 3.0 Example Server Architecture

[0056] FIG. **4** is a block diagram of a server **200**, which may be used in the cloud-based system **100**, in other systems, or standalone. For example, the enforcement nodes **150** and the central

authority **152** may be formed as one or more of the servers **200**. The server **200** may be a digital computer that, in terms of hardware architecture, generally includes a processor **202**, input/output (I/O) interfaces **204**, a network interface **206**, a data store **208**, and memory **210**. It should be appreciated by those of ordinary skill in the art that FIG. **4** depicts the server **200** in an oversimplified manner, and a practical embodiment may include additional components and suitably configured processing logic to support known or conventional operating features that are not described in detail herein. The components (**202**, **204**, **206**, **208**, and **210**) are communicatively coupled via a local interface **212**. The local interface **212** may be, for example, but not limited to, one or more buses or other wired or wireless connections, as is known in the art. The local interface **212** may have additional elements, which are omitted for simplicity, such as controllers, buffers (caches), drivers, repeaters, and receivers, among many others, to enable communications. Further, the local interface **212** may include address, control, and/or data connections to enable appropriate communications among the aforementioned components.

[0057] The processor **202** is a hardware device for executing software instructions. The processor **202** may be any custom made or commercially available processor, a Central Processing Unit (CPU), an auxiliary processor among several processors associated with the server **200**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the server **200** is in operation, the processor **202** is configured to execute software stored within the memory **210**, to communicate data to and from the memory **210**, and to generally control operations of the server **200** pursuant to the software instructions. The I/O interfaces **204** may be used to receive user input from and/or for providing system output to one or more devices or components.

[0058] The network interface **206** may be used to enable the server **200** to communicate on a network, such as the Internet **104**. The network interface **206** may include, for example, an Ethernet card or adapter or a Wireless Local Area Network (WLAN) card or adapter. The network interface **206** may include address, control, and/or data connections to enable appropriate communications on the network. A data store **208** may be used to store data. The data store **208** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, and the like), and combinations thereof.

[0059] Moreover, the data store **208** may incorporate electronic, magnetic, optical, and/or other types of storage media. In one example, the data store **208** may be located internal to the server **200**, such as, for example, an internal hard drive connected to the local interface **212** in the server **200**. Additionally, in another embodiment, the data store **208** may be located external to the server **200** such as, for example, an external hard drive connected to the I/O interfaces **204** (e.g., SCSI or USB connection). In a further embodiment, the data store **208** may be connected to the server **200** through a network, such as, for example, a network-attached file server.

[0060] The memory **210** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, etc.), and combinations thereof. Moreover, the memory **210** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **210** may have a distributed architecture, where various components are situated remotely from one another but can be accessed by the processor **202**. The software in memory **210** may include one or more software programs, each of which includes an ordered listing of executable instructions for implementing logical functions. The software in the memory **210** includes a suitable Operating System (O/S) **214** and one or more programs **216**. The operating system **214** essentially controls the execution of other computer programs, such as the one or more programs **216**, and provides scheduling, input-output control, file and data management, memory management, and communication control and related services. The one or more programs **216** may be configured to implement the various processes, algorithms, methods, techniques, etc. described

herein.

#### § 4.0 Example User Device Architecture

[0061] FIG. 5 is a block diagram of a user device **300**, which may be used with the cloud-based system **100** or the like. Specifically, the user device **300** can form a device used by one of the users **102**, and this may include common devices such as laptops, smartphones, tablets, netbooks, personal digital assistants, MP3 players, cell phones, e-book readers, IoT devices, servers, desktops, printers, televisions, streaming media devices, and the like. The user device **300** can be a digital device that, in terms of hardware architecture, generally includes a processor **302**, I/O interfaces **304**, a network interface **306**, a data store **308**, and memory **310**. It should be appreciated by those of ordinary skill in the art that FIG. 5 depicts the user device **300** in an oversimplified manner, and a practical embodiment may include additional components and suitably configured processing logic to support known or conventional operating features that are not described in detail herein. The components (**302**, **304**, **306**, **308**, and **302**) are communicatively coupled via a local interface **312**. The local interface **312** can be, for example, but not limited to, one or more buses or other wired or wireless connections, as is known in the art. The local interface **312** can have additional elements, which are omitted for simplicity, such as controllers, buffers (caches), drivers, repeaters, and receivers, among many others, to enable communications. Further, the local interface **312** may include address, control, and/or data connections to enable appropriate communications among the aforementioned components.

[0062] The processor **302** is a hardware device for executing software instructions. The processor **302** can be any custom made or commercially available processor, a CPU, an auxiliary processor among several processors associated with the user device **300**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the user device **300** is in operation, the processor **302** is configured to execute software stored within the memory **310**, to communicate data to and from the memory **310**, and to generally control operations of the user device **300** pursuant to the software instructions. In an embodiment, the processor **302** may include a mobile optimized processor such as optimized for power consumption and mobile applications. The I/O interfaces **304** can be used to receive user input from and/or for providing system output. User input can be provided via, for example, a keypad, a touch screen, a scroll ball, a scroll bar, buttons, a barcode scanner, and the like. System output can be provided via a display device such as a Liquid Crystal Display (LCD), touch screen, and the like.

[0063] The network interface **306** enables wireless communication to an external access device or network. Any number of suitable wireless data communication protocols, techniques, or methodologies can be supported by the network interface **306**, including any protocols for wireless communication. The data store **308** may be used to store data. The data store **308** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDRom, and the like), and combinations thereof. Moreover, the data store **308** may incorporate electronic, magnetic, optical, and/or other types of storage media.

[0064] The memory **310** may include any of volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, etc.), and combinations thereof. Moreover, the memory **310** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **310** may have a distributed architecture, where various components are situated remotely from one another but can be accessed by the processor **302**. The software in memory **310** can include one or more software programs, each of which includes an ordered listing of executable instructions for implementing logical functions. In the example of FIG. 3, the software in the memory **310** includes a suitable operating system **314** and programs **316**. The operating system **314** essentially controls the execution of other computer programs and provides scheduling, input-output control, file and

data management, memory management, and communication control and related services. The programs **316** may include various applications, add-ons, etc. configured to provide end user functionality with the user device **300**. For example, example programs **316** may include, but not limited to, a web browser, social networking applications, streaming media applications, games, mapping and location applications, electronic mail applications, financial applications, and the like. In a typical example, the end-user typically uses one or more of the programs **316** along with a network such as the cloud-based system **100**.

#### § 5.0 Zero Trust Network Access Using the Cloud-Based System

[0065] FIG. **6** is a network diagram of a zero-trust network access (ZTNA) application utilizing the cloud-based system **100**. For ZTNA, the cloud-based system **100** can dynamically create a connection through a secure tunnel between an endpoint (e.g., users **102A**, **102B**) that are remote and an on-premises connector **400** that is either located in cloud file shares and applications **402** and/or in an enterprise network **410** that includes enterprise file shares and applications **404**. The connection between the cloud-based system **100** and on-premises connector **400** is dynamic, on-demand, and orchestrated by the cloud-based system **100**. A key feature is its security at the edge—there is no need to punch any holes in the existing on-premises firewall. The connector **400** inside the enterprise (on-premises) “dials out” and connects to the cloud-based system **100** as if too were an endpoint. This on-demand dial-out capability and tunneling authenticated traffic back to the enterprise is a key differentiator for ZTNA. Also, this functionality can be implemented in part by the application **350** on the user device **300**. Also, the applications **402**, **404** can include B2B applications. Note, the difference between the applications **402**, **404** is the applications **402** are hosted in the cloud, whereas the applications **404** are hosted on the enterprise network **410**. The B2B service described herein contemplates use with either or both of the applications **402**, **404**.

[0066] The paradigm of virtual private access systems and methods is to give users network access to get to an application and/or file share, not to the entire network. If a user is not authorized to get the application, the user should not be able even to see that it exists, much less access it. The virtual private access systems and methods provide an approach to deliver secure access by decoupling applications **402**, **404** from the network, instead of providing access with a connector **400**, in front of the applications **402**, **404**, an application on the user device **300**, a central authority **152** to push policy, and the cloud-based system **100** to stitch the applications **402**, **404** and the software connectors **400** together, on a per-user, per-application basis.

[0067] With the virtual private access, users can only see the specific applications **402**, **404** allowed by the central authority **152**. Everything else is “invisible” or “dark” to them. Because the virtual private access separates the application from the network, the physical location of the application **402**, **404** becomes irrelevant-if applications **402**, **404** are located in more than one place, the user is automatically directed to the instance that will give them the best performance. The virtual private access also dramatically reduces configuration complexity, such as policies/firewalls in the data centers. Enterprises can, for example, move applications to Amazon Web Services or Microsoft Azure, and take advantage of the elasticity of the cloud, making private, internal applications behave just like the marketing leading enterprise applications. Advantageously, there is no hardware to buy or deploy because the virtual private access is a service offering to end-users and enterprises.

#### § 6.0 Digital Experience Monitoring

[0068] FIG. **7** is a network diagram of the cloud-based system **100** in an application of digital experience monitoring. Here, the cloud-based system **100** providing security as a service as well as ZTNA, can also be used to provide real-time, continuous digital experience monitoring, as opposed to conventional approaches (synthetic probes). A key aspect of the architecture of the cloud-based system **100** is the inline monitoring. This means data is accessible in real-time for individual users from end-to-end. As described herein, digital experience monitoring can include monitoring, analyzing, and improving the digital user experience.

[0069] The cloud-based system **100** connects users **102** at the locations **110, 112, 118** to the applications **402, 404**, the Internet **104**, the cloud services **106**, etc. The inline, end-to-end visibility of all users enables digital experience monitoring. The cloud-based system **100** can monitor, diagnose, generate alerts, and perform remedial actions with respect to network endpoints, network components, network links, etc. The network endpoints can include servers, virtual machines, containers, storage systems, or anything with an IP address, including the Internet of Things (IoT), cloud, and wireless endpoints. With these components, these network endpoints can be monitored directly in combination with a network perspective. Thus, the cloud-based system **100** provides a unique architecture that can enable digital experience monitoring, network application monitoring, infrastructure component interactions, etc. Of note, these various monitoring aspects require no additional components—the cloud-based system **100** leverages the existing infrastructure to provide this service.

[0070] Again, digital experience monitoring includes the capture of data about how end-to-end application availability, latency, and quality appear to the end user from a network perspective. This is limited to the network traffic visibility and not within components, such as what application performance monitoring can accomplish. Networked application monitoring provides the speed and overall quality of networked application delivery to the user in support of key business activities. Infrastructure component interactions include a focus on infrastructure components as they interact via the network, as well as the network delivery of services or applications. This includes the ability to provide network path analytics.

[0071] The cloud-based system **100** can enable real-time performance and behaviors for troubleshooting in the current state of the environment, historical performance and behaviors to understand what occurred or what is trending over time, predictive behaviors by leveraging analytics technologies to distill and create actionable items from the large dataset collected across the various data sources, and the like. The cloud-based system **100** includes the ability to directly ingest any of the following data sources network device-generated health data, network device-generated traffic data, including flow-based data sources inclusive of NetFlow and IPFIX, raw network packet analysis to identify application types and performance characteristics, HTTP request metrics, etc. The cloud-based system **100** can operate at 10 gigabits (10G) Ethernet and higher at full line rate and support a rate of 100,000 or more flows per second or higher.

[0072] The applications **402, 404** can include enterprise applications, Office 365, Salesforce, Skype, Google apps, internal applications, etc. These are critical business applications where user experience is important. The objective here is to collect various data points so that user experience can be quantified for a particular user, at a particular time, for purposes of analyzing the experience as well as improving the experience. In an embodiment, the monitored data can be from different categories, including application-related, network-related, device-related (also can be referred to as endpoint-related), protocol-related, etc. Data can be collected at the application **350** or the cloud edge to quantify user experience for specific applications, i.e., the application-related and device-related data. The cloud-based system **100** can further collect the network-related and the protocol-related data (e.g., domain name system (DNS) response time).

Application-Related Data

TABLE-US-00001 Page Load Time Redirect count (#) Page Response Time Throughput (bps) Document Object Model Total size (bytes) (DOM) Load Time Total Downloaded bytes Page error count (#) App availability (%) Page element count by category (#)

Network-Related Data

TABLE-US-00002 HTTP Request metrics Bandwidth Server response time Jitter Ping packet loss (%) Trace Route Ping round trip DNS lookup trace Packet loss (%) GRE/IPSec tunnel monitoring Latency MTU and bandwidth measurements

Device-Related Data (Endpoint-Related Data)

TABLE-US-00003 System details Network (config) Central Processing Unit (CPU) Disk Memory

(RAM) Processes Network (interfaces) Applications

[0073] Metrics could be combined. For example, device health can be based on a combination of CPU, memory, etc. Network health could be a combination of Wi-Fi/LAN connection health, latency, etc. Application health could be a combination of response time, page loads, etc. The cloud-based system **100** can generate service health as a combination of CPU, memory, and the load time of the service while processing a user's request. The network health could be based on the number of network path(s), latency, packet loss, etc.

[0074] The lightweight connector **400** can also generate similar metrics for the applications **402**, **404**. In an embodiment, the metrics can be collected while a user is accessing specific applications that user experience is desired for monitoring. In another embodiment, the metrics can be enriched by triggering synthetic measurements in the context of an in-line transaction by the application **350** or cloud edge. The metrics can be tagged with metadata (user, time, app, etc.) and sent to a logging and analytics service for aggregation, analysis, and reporting. Further, network administrators can get UEX reports from the cloud-based system **100**. Due to the inline nature and the fact the cloud-based system **100** is an overlay (in-between users and services/applications), the cloud-based system **100** enables the ability to capture user experience metric data continuously and to log such data historically. As such, a network administrator can have a long-term detailed view of the network and associated user experience.

## § 7.0 SSL Overview

[0075] Secure Sockets Layer (SSL) is a client-server protocol that creates a secure channel over the Internet. SSL is used to validate the identity of the destination server and (optionally) the client, and to encrypt information sent across the internet between the client and server. FIG. **8** is a flow diagram illustrating an example handshake for HTTPS to describe a secure, encrypted tunnel between a client (e.g., the user device **300**) and a server (e.g., the server **200**). When a client, such as a browser, first sends an HTTPS request to a server, it starts a series of message exchanges called the SSL handshake. The client can send an HTTPS request with supported cipher suites and compression algorithms, session ID, SSL version, and a randomly generated value, i.e., a “client hello.”

[0076] The server sends its digital certificate to the client to authenticate itself, as well as the selected cipher suite and compression algorithm, session ID, SSL session, a randomly generated value, a certificate with a public key, and optionally a request for the client's certificate, i.e., a “server hello.” The client verifies the certificate with a Certificate Authority (CA), sends the pre-master secret computed with both random values, and encrypted with the server's public key. The client notifies the server that all subsequent messages will be encrypted with the keys and negotiated algorithms, i.e., the client and server agree on the SSL protocol version and algorithms to use, and the client and server generate the symmetric keys they will use to encrypt their messages.

[0077] The server uses its private key to decrypt the pre-master key, only the server with the private key that matches the public key that was sent with the certificate can decrypt the pre-master key. The server validates the browser (client) certificate and uses the public key to decrypt the messages. The server notifies the client that all subsequent messages will be encrypted using the keys and negotiated algorithms. The server computes the master key from the pre-master key and generates the session key. The server sends a message that is a hash of the exchanged messages using the master key and the session key. The client decrypts the message and validates the hash, leading to a successful handshake.

[0078] After the SSL handshake is successfully completed, the client and server continue with the standard HTTP communications in a secure manner.

[0079] FIG. **9** is a screenshot of packet capture showing SSL packets as they are exchanged between a client and a server. The client sends its HTTPS request in the Client Hello. The entire HTTPS message is encrypted, including the headers and the request/response load. The actual

hostname and domain name being accessed is not visible. How the cloud-based system **100** determines the destination hostname depends on whether it is operating in transparent mode or explicit mode. The server responds with its Hello message and its certificate. (A certificate is an electronic form that verifies the identity and public key of the subject of the certificate.) SSL uses the Public Key Infrastructure (PKI) to ensure the trustworthiness of the certificates. The client and server continue with the SSL negotiation. After the SSL tunnel is established, the application data is sent securely through the tunnel.

[0080] SSL uses Public Key Infrastructure (PKI) to ensure the trustworthiness of the certificates. PKI uses a trusted third party, called a Certificate Authority (CA), to guarantee the identity of an entity. When a CA verifies an entity's identity, it uses an algorithm, such as RSA, to generate a public and private key. It gives the private key to the requesting entity, and the public key is made available to the public. To authenticate itself to another party, the entity uses its private key to encrypt its certificate, and the other party uses the corresponding public key to decrypt it.

[0081] A CA issues certificates in a tree structure, with the root certificate as the top-most certificate. The CA signs the root certificate, which is considered trustworthy in many software applications, such as web browsers. Web browsers have the root certificates of many CAs.

[0082] A root certificate can sign and designate a certificate as an intermediate CA certificate, which can sign and designate other certificates as intermediate certificates as well. A certificate chain refers to the list of certificates that complete the chain of trust, from the trusted root CA certificate to any intermediate certificates and the certificate of an entity. The following is an example of a certificate chain.

TABLE-US-00004 The certificate of mail.google.com was signed by Google Internet Authority G2. The certificate of Google Internet Authority G2 was signed by GeoTrust Global CA. The certificate of GeoTrust Global was signed by Equifax Secure Certificate Authority. The certificate of GeoTrust Global CA and Equifax Secure Certificate Authority are in the certificate store of the browser.

### Perfect Forward Secrecy (PFS)

[0083] Perfect Forward Secrecy (PFS) is a feature of secure communication protocols that prevent compromised session keys. In the commonly used RSA key exchange, SSL sessions between the client and web server are encrypted with the public key and decrypted with the private key. If attackers access the server's private key, they can uncover the session keys and decrypt all conversations from past and future sessions.

[0084] In contrast, PFS uses either the standard Diffie-Hellman ephemeral key exchange (DHE) or the Elliptic Curve Diffie-Hellman ephemeral key exchange (ECDHE). DHE uses public-key cryptography, which generates keys with modular arithmetic. In DHE, there is not a link between the server's private key and session key, so the confidentiality of session keys are not dependent on the private keys. If attackers access the server's private key, they are unable to uncover the session key and decrypt the conversation. Furthermore, the server generates different session keys for each conversation with the client. If attackers compromise the session key, they are only able to decrypt the conversation for that particular session. To decrypt all conversations, they must compromise the session keys for every session.

[0085] ECDHE is like DHE but uses elliptic-curve cryptography. Elliptic-curve cryptography generates keys using algebraic curves. It is significantly faster than DHE and provides better performance. Elliptic-curve cryptography achieves equivalent security as RSA with smaller keys.

### § 7.1 SSL Inspection

[0086] HTTPS is an aggregate of HTTP and the SSL/TLS protocol, wherein the authentication and encryption capabilities of SSL/TLS protect HTTP communications. This is vital because the information that is sent on the Internet is passed along from one device to another before it reaches the destination server. Therefore, sensitive information, such as credit card numbers, usernames, and passwords, may be seen by intermediate devices if the information is sent in clear text over

HTTP. When the information is encrypted and protected by the SSL protocol, only the intended recipient can read the information.

[0087] Unfortunately, the security provided by SSL is also being misused in a number of ways:

[0088] SSL encryption is used to hide dangerous content such as viruses, spyware, and other malware.

[0089] Attackers build their websites with SSL encryption.

[0090] Attackers inject their malicious content into well-known and trusted SSL-enabled sites.

[0091] SSL can be used to hide data leakage, for example, the transmission of sensitive financial documents from an organization or the like.

[0092] SSL can be used to hide the browsing of websites that belong to legal-liability classes.

[0093] As more and more websites use HTTPS, including social media, the ability to control and inspect traffic to and from these sites has become an important piece of the security posture of an organization.

[0094] The cloud-based system **100** can inspect HTTPS traffic from an organization. The service can scan data transactions and apply policies to it, as described herein. An enforcement node **150** can function as a full SSL proxy, or SSL man-in-the-middle (MITM) proxy.

[0095] The cloud-based system **100** can provide two options to protect HTTPS traffic: SSL inspection, or if SSL inspection is not feasible, one can configure a global block of specific HTTPS content.

[0096] FIG. **10** is a flow diagram illustrating an embodiment of SSL inspection **450** with the cloud-based system **100** as a proxy. In this embodiment, the cloud-based system **100** establishes a separate SSL tunnel with the user's browser and with the destination server. FIG. **10** illustrates the SSL inspection **450** process. First, a user (at the user device **300**) opens a browser and sends an HTTPS request. Second, the cloud-based system **100** intercepts the HTTPS request. Through a separate SSL tunnel, the cloud-based system **100** sends its HTTPS request to the destination server (the server **200**) and conducts SSL negotiations. The destination server sends the cloud-based system **100** its certificate with its public key. The cloud-based system **100** and destination server complete the SSL handshake. The application data and subsequent messages are sent through the SSL tunnel. The cloud-based system **100** conducts SSL negotiations with the user's browser. It sends the browser an intermediate certificate or an organization's custom intermediate root as well as a server certificate signed by the intermediate CA. The browser validates the certificate chain in the browser's certificate store. The cloud-based system **100** and the browser complete the SSL handshake. The application data and subsequent messages are sent through the SSL tunnel.

[0097] In an embodiment, the SSL inspection can use an intermediate certificate of the cloud-based system **100**. With this option, the cloud-based system **100** dynamically generates and signs the server certificate that it presents to the client. This certificate contains the same fields as the original destination server certificate, except for the identifying information of the issuer, called the issuer distinguished name (DN). The issuer DN is set to the name of the cloud-based system **100** intermediate certificate. The browser receives this certificate signed by the cloud-based system **100** intermediate certificates along with the cloud-based system **100** intermediate certificate. To enable a browser or system to automatically trust all certificates signed by the cloud-based system **100** Certificate Authority, users must install the cloud-based system **100** Root CA certificate on their workstations.

[0098] In another embodiment, the SSL inspection can use a custom intermediate root certificate. One can subscribe to the Custom Certificate feature and configure a custom intermediate root certificate for SSL inspection. Here, the cloud-based system **100** does not use an organization's root certificate or private keys. Instead, it uses the custom intermediate root certificate signed by a trusted CA, so it is possible to use a CA that is already deployed on an organization's machines. To configure an intermediate root certificate, the cloud-based system **100** generates a certificate signing request (CSR) with a key pair (i.e., public and private key) and encrypts the private key



using AES. The private key is stored securely in the central authority **152**, while the CSR contains the public key.

[0099] After the CA signs the CSR, the signed certificate can be uploaded to the cloud-based system **100**. During the SSL negotiation with the user's browser, the cloud-based system **100** dynamically generates and signs the server certificate that it presents to the client with this intermediate certificate. The certificate issuer is set to the organization name, and the cloud-based system **100** generates the certificate once per site and caches these certificates on the enforcement node **150**. These cached certificates are usually valid until their expiration date.

[0100] In addition to the intermediate root certificate, it is possible to upload the certificate chain that includes any other intermediate certificates that complete the chain to the intermediate root certificate. When the certificate chain is uploaded, the cloud-based system **100** sends the intermediate root certificate along with this key chain and the signed server certificate to the users' machines during SSL inspection. If the certificate chain is not uploaded, the cloud-based system **100** sends only the organization's intermediate root certificate and its signed server certificate to the user's machine. Uploading the certificate chain provides important benefits. The certificate chain ensures that the users' machines can validate the server certificate signed by the organization's intermediate CA even if the users' browsers have only the root certificate in their certificate store. If the certificate is changed due to the compromise of an intermediate root certificate, or simply as a routine security measure, the ability to send the certificate chain to users' machines during SSL inspection is a key benefit. Because it enables certificate rotation efficiently without the need for a new key ceremony or certificate push to an organization's users.

[0101] The cloud-based system **100** provides a CRL (Certificate Revocation List) distribution point (CDP) for every certificate it generates so that client applications can locate the Certificate Revocation Lists (ORLs) as necessary.

## § 7.2 SSL Handshake Process

[0102] FIG. **11** is a flow diagram of details of an SSL handshake process **500** between an SSL client **502** and an SSL server **504**. The SSL client **502** can be the user device **300**, etc. and the SSL server **504** can be a location on the Internet **104**, etc., i.e., the server **200**. That is, the SSL server **504** can be an endpoint for an encrypted tunnel with the user device **300**. The SSL client **502** sends a “client hello” message that lists cryptographic information such as the SSL version and, in the client's order of preference, the CipherSuites supported by the SSL client **502** (step **510-1**). The message also contains a random byte string that is used in subsequent computations. The protocol allows for the “client hello” to include the data compression methods supported by the SSL client **502**.

[0103] The SSL server **504** responds with a “server hello” message that contains the CipherSuite chosen by the SSL server **504** from the list provided by the SSL client **502**, the session ID, and another random byte string (step **510-2**). The SSL server **504** also sends its digital certificate. If the SSL server **504** requires a digital certificate for client authentication, the SSL server **504** sends a “client certificate request” that includes a list of the types of certificates supported and the Distinguished Names of acceptable CAs. The SSL client **502** verifies the SSL server's **504** digital certificate (step **510-3**).

[0104] The SSL client **502** sends the random byte string that enables both the SSL client **402** and the SSL server **504** to compute the secret key to be used for encrypting subsequent message data (step **510-4**). The random byte string itself is encrypted with the SSL server's **404** public key. If the SSL server **504** sent a “client certificate request,” the SSL client **502** sends a random byte string encrypted with the client's private key, together with the SSL client's **502** digital certificate, or a “no digital certificate alert” (step **510-5**). This alert is only a warning, but with some implementations, the handshake fails if client authentication is mandatory. The SSL server **504** verifies the client's certificate if required (step **510-6**).

[0105] The SSL client **502** sends the server a “finished” message, which is encrypted with the

secret key, indicating that the SSL client **502** part of the handshake is complete (step **510-7**). The SSL server **504** sends the SSL client **502** a “finished” message, which is encrypted with the secret key, indicating that the SSL server **504** part of the handshake is complete. For the duration of the SSL session, the SSL server **504** and SSL client **502** can now exchange messages that are symmetrically encrypted with the shared secret key (step **510-9**).

### § 7.3 SSL Interception Proxies

[0106] FIG. **12** is a flow diagram of a process **520** performing SSL interception through an interception proxy **530** in the handshake process **500**. The interception proxy **530** can be one of the enforcement nodes **150** in the cloud-based system **100**. Enterprises deploy or use the interception proxy **530** to secure themselves from SSL-based threats, which are increasingly common. The interception proxy **530** works by acting as a MITM and modifying the encrypted channel. Whenever the SSL client **502** initiates a connection to a remote SSL server **504**, the interception proxy **530** will intercept it and open two different channels of communication, one with the SSL client **502** and the other with the SSL server **504** that the SSL client **502** intended to talk to in the first place. This allows the interception proxy **530** to actively modify/inject the content from the SSL client **502** to the SSL server **504** or vice versa. This allows IT admins to perform malware scanning and other security functions on the otherwise encrypted content. In order to achieve this, an IT admin usually deploys proxy's ROOT CA certificate on the user devices **300** for the SSL clients **502** to trust the handshake which happens between the SSL client **502** and the interception proxy **530** which generates a certificate for every SSL server **504** that the SSL client **502** tries to communicate with. This naturally breaks with apps that employ certificate pinning for enhanced security. Note, while the term “interception proxy” is used herein, those skilled in the art will recognize this is a functional name and it may be called other names while being the same based on the functionality.

[0107] Advantageously, the interception proxy **530** enables interception, inspection, and filtering of content on an otherwise encrypted channel. For example, the cloud-based system **100** using the interception proxy **530** can perform DLP, web content filtering, malware detection, intrusion detection/prevention, firewall and Deep Packet Inspection (DPI), etc. The interception proxy **530** acts as the SSL client **502** on the SSL server **504** side and as the SSL server **504** on the SSL client **502** sides.

[0108] The interception proxy **530** performs SSL inspection by breaking or terminating the encrypted tunnel in the cloud-based system **100**. Specifically, the enforcement node **150** is a proxy, and it has an encrypted tunnel with the client and another encrypted tunnel with the server. That is, this approach requires SSL/TLS/DTLS handshake/termination on the enforcement node **150** (in the cloud, on-premises, etc.). This approach, with the enforcement node **150** as a MITM proxy breaking the tunnel has limitations. Specifically, some applications use Certificate Pinning or other techniques to prevent MITM. With Certificate Pinning, the client is configured to only accept a specific certificate or a specific CA. In this case, the application will break when presented with a certificate signed by the cloud-based system **100**, even if it is trusted.

[0109] This is done to ensure greater control over the communicating entities and to prevent the MITM attacks. The situation is somewhat of a paradox: entities such as Domain Name Systems (DNS) and CAs are trusted and supposed to supply trusted input. However, more and more applications are trying hard with pinning to eliminate this conference of trust. By pinning the certificate or the public key of the server certificate, an application no longer needs to depend on third-party entities such as DNS, CA, etc. when making security decisions relating to a peer's identity. This makes an app immune to MITM attacks. Pinning effectively removes the “conference of trust” by eliminating the set of entities that are beyond the control of a domain owner. Apps achieve this by accepting server certificates that strictly match a defined criterion, usually subject key information.

[0110] With the SSL interception, proxy servers are employed in the cloud-based system **100** are

aware of the SSL encrypted communication and may need to intercept it in order to provide security services. Such filtering solutions are generally achieved through interception proxies that engage in deep packet inspection to resist SSL-based threats that may range from trivial viruses to sophisticated ransomware. The problem when apps employ certificate pinning is that they reject the connection during negotiation with an interception proxy on account of peer's (in this case, SSL proxy) untrusted certificate.

[0111] Such apps fail to function in the enterprise environment and fail to provide desired services leading to bad user experience and frustration. The apps would be rendered dysfunctional partially or completely due to the certificate pinning employed by them. They will terminate the connection upon receiving a server certificate from the proxy that does not match the criterion. This leads to bad user experience, and the cloud security system does not have any visibility or resolution of such issues.

[0112] As more and more viruses use encrypted channels to infect machines, it is imperative for enterprises to employ SSL interception proxies to protect users. This poses a conundrum as app developers would like to eliminate trust on third parties like CAs, which may be vulnerable to other attacks. To solve this issue, an IT admin may be lured to turn SSL interception off, which makes their enterprise security even worse. Hence, it is desirable for IT admins to selectively turn SSL interception off only for some trusted applications and domains. Since it is very hard for IT admins to know apriori which apps users will use or what domains the app may hit, which may even change over time, there is a huge need for a better tunneling solution.

[0113] The cloud-based system **100** has little or no idea about the dysfunctional apps. The client apps terminate the connection with or without an alert message to the server upon receiving the mismatched certificate. Further, the IT admin has no way to find all the apps and their server domains for which the app performs pinning. As a result, this design does not allow the users to use such apps while subscribing to the security or enterprise compliance policies. To make these apps functional again, the cloud-based system **100** cannot perform the SSL interception described in FIG. **10**, e.g., bypass SSL interception.

#### § 7.4 SSL Interception

[0114] FIG. **13** is a network diagram of a network **600** with the enforcement node **150** configured as an interception proxy **530**. As such, an interception proxy **530** in the cloud-based system **100** can selectively intercept SSL communications. In an embodiment, Internet-bound traffic of the user device **300** (the SSL client **502**) is controlled through a tunnel **610** to the cloud-based system **100** which has a second tunnel **612** to the SSL server **504**. The tunnel **610** acts as an intermediary passive MITM proxy that relays all the network requests and responses from client applications **620** to the cloud-based system **100**. To achieve this, a process running on the host (the SSL client **502**) installs a virtual interface on the user device **300**. The process installs a default route on the interface in the device routing table and opens listening sockets for User Datagram Protocol (UDP) and Transmission Control Protocol (TCP) traffic at randomly available ports.

#### § 7.5 SSL Inspection Based on Key Snooping

[0115] FIG. **14** is a network diagram of a network **650** with the enforcement node **150** operating as a snooping proxy **660** to perform SSL interception without breaking the tunnel as with the interception proxy **530**. This presents a different approach for SSL interception than the interception proxy **530**, which avoids the disadvantages of certificate pinning and certificate management. In the network **650**, a tunnel **670** is between the SSL client **502** and the SSL server **504**. Again, the tunnel **670** can be SSL, TLS, DTLS, HTTPS, etc. The key difference with the snooping proxy **660** relative to the interception proxy **530** is the snooping proxy **660** does not break the tunnel **670**. Note, the snooping proxy **660** is still a MITM proxy like the interception proxy **530**. Again, while the term “snooping proxy” is used herein, those skilled in the art will recognize this is a functional name and it may be called other names while being the same based on the functionality.

[0116] The snooping proxy **660** can be one of the enforcement nodes **150** in the cloud-based system **100**. Also, the client **502** can be the user device **300** including the application **350**. As described herein, the application **350** is a traffic-forwarding application that enables the user device **300** to operate (communicate) with the cloud-based system **100**. The snooping proxy **660**, being already a MITM proxy, can snoop (monitor) on the handshake process **500**. This snooping can be at the enforcement node **150** operating as the snooping proxy **660** as well as at the application **350**. This snooping can also use key agents, such as part of the application **350**, operating system support hooks, such as at the user device **300**, etc. The key aspect here is the snooping proxy **660** can snoop the handshake process **500** for purposes of obtaining keys.

[0117] Once the snooping proxy **660** has keys for a given session, the snooping proxy **660** can monitor the encrypted traffic on the tunnel **670**. Note, typically, monitoring in the cloud-based system **100** is inline in a sense the enforcement node **150** sits directly between the client **502** (the user device **300**) and the server **504** (or any other destination on the Internet **104**, the cloud services **106**, etc.). Here, the snooping proxy **660** is still inline. The snooping proxy **660** can receive encrypted traffic, view and inspect the traffic based on the snooping of the keys, and allow or block the traffic based on the inspection.

[0118] This approach solves the various limitations with a traditional MITM proxy as an interception proxy **530**. That is, applications with certificate pinning now can support SSL inspection to block policy violations or malware transfers. This removes the need for certificate deployments with the cloud-based system **100**. Also, it is possible to decode any other variant of SSL to inspect or detect application signature (aka DPI) inside an encapsulated layer or protocol. Further, this approach is completely transparent to primitive SSL-based applications such as FTPS, which cannot trust MITM root certificates. Finally, this allows granular policy control and transactional visibility for critical or productivity applications without breaking the SSL protocol.

#### § 7.6 SSL Profile Construction, Learning and Transfer of Knowledge

[0119] In either SSL environment, namely the interception proxy **530** and the snooping proxy **660**, for every new connection, the application **350** process on the device can create a state machine or the like for the transaction, and, based on the results of the transaction, the process constructs a profile for the SSL client **502** which initiated the connection. For every connection, the process can construct a profile for the connection as a tuple: <Origin, Host-Name, Destination-Socket-Address, Handshake-Status, Key information>.

[0120] The origin is the client application **620**, which is originating a request. The origin information is obtained through a process to port mapping on the host machine. The Host Name is the fully qualified domain name of the SSL server **504** that the SSL client **502** is trying to reach. The hostname is retrieved from the SNI (Server Name Indication) parsed as a TLS extension in the Client Hello SSL record. The Destination Socket contains information about Destination-Server-IP-Address: Destination-Port that the SSL client **502** is trying to establish a connection. This information is retrieved by parsing the IP-packet header during connection establishment.

[0121] The Handshake Status is a bit flag that keeps a record of SSL handshake messages exchanged with the SSL server **504**. The flag is set to 1 if the handshake succeeds, and the client starts sending Application Data to the server. The profile is learned for every transaction and reevaluated whenever the SSL client **502** tries to reach the same destination. This knowledge is periodically transferred to the cloud-based system **100** out-of-band on a persistent control channel that allows the cloud-based system **100** to learn the behavior of client apps **620** with SSL interception.

[0122] To construct this profile, the process passively observes the SSL Record Layer data messages and keep track of all the records that have been exchanged for any given transaction. For example, the process can parse the SSL headers to check if the SSL client **502** returns an SSL alert and/or if application data is sent over the connection. The process can parse the initial (K) server bytes and check the intermediate CA certificate from the enforcement node **150**. The process can

find the processes and host corresponding to the connection.

[0123] The following SSL handshake messages can be recorded:

[0124] Client hello to determine the SSL server **504** the SSL client **502** wants to connect with. The SNI host field provides the information.

[0125] Server Hello to determine the server response towards the client request and client supported ciphers.

[0126] A certificate that contains the certificates advertised by the SSL server **504** and which is used to check if SSL interception is enabled for the transaction.

[0127] Alert (optional), which indicates if the SSL client **502** rejected the certificate and the reason for rejection.

[0128] Application data which indicates the successful handshake since the application data is exchanged now.

[0129] This process can be extended to generate more detailed profiles containing the ciphers supported by the SSL client **502** and the SSL server **504**, SSL version, certificate chain, etc.

[0130] Every SSL message is sent as part of the Record Layer Protocol which provides messages in the following format:

TABLE-US-00005 Content type (1 Octet) Version (2 Octets) Length (2 Octets) Data

§ 7.7 Security Functions on Traffic with SSL Inspection, Either with the Interception Proxy or the Snooping Proxy

[0131] The cloud-based system **100** can support various security functions on encrypted traffic, including:

[0132] Granular URL filtering and cloud app control policies where the cloud-based system **100** can enforce granular user, group, and location policies that not only control access to sites or applications but also control what a user can do within an application. For example, it is possible to define a Web email policy that allows users to view and send mail, but not attachments, or a social media policy that allows users to view Facebook, but not post.

[0133] Skipping Inspection for Specific URLs/URL categories: When configuring SSL Inspection policy, it is possible to prevent the service from inspecting sessions to certain URLs or URL categories (for example, in the Banking and Healthcare URL categories). This list can apply globally through an organization as well as granular to users, groups of users, etc.

[0134] Skipping Inspection for Specific Cloud Applications/Cloud Application Categories: When configuring SSL Inspection policy, it is possible to prevent the cloud-based system **100** from inspecting transactions to specific cloud applications or cloud application categories. This list can apply globally through an organization as well as granular to users, groups of users, etc.

[0135] Content Filtering where the cloud-based system **100** is enabled to block malicious or inappropriate content in a page, such as during a Google search.

[0136] Block Undecryptable Transactions: wherein the cloud-based system **100** is configured to block the transactions of applications that the cloud-based system **100** cannot decrypt because of using non-standard encryption methods and algorithms, as well as where snooping fails and where the interception proxy **530** encounters certificate pinning.

[0137] Block Advanced Persistent Threats (APT) in encrypted traffic. Note, most targeted malware is now delivered over SSL.

[0138] Control access to Google consumer apps and non-corporate Google accounts.

[0139] Block access to sites with revoked certificates: The cloud-based system **100** supports OCSP (Online Certificate Status Protocol) to verify the validity of all server certificates. It verifies the OCSP responder URL in a server's certificate and sends an OCSP request to the responder. The cloud-based system **100** allows access if the responder indicates that the certificate is Good, and blocks access if the responder responds that the certificate is Unknown or Revoked. The cloud-based system **100** displays a notification when it blocks access to a site due to a bad certificate (if the certificate issuer is unknown, or if the certificate has expired, or if the Common Name in the

certificate does not match). It also logs these transactions with “bad server cert” in the policy field. [0140] Data Loss Prevention (DLP): The cloud-based system **100** can enforce the DLP policy when SSL inspection is enabled.

[0141] Of note, the enforcement node **150** can be configured, not as a caching proxy. Data is inspected in the enforcement node's **150** memory after decryption and sent out to the client immediately. Even when a core dump is taken on the enforcement node **150**, SSL (encrypted) session data is cleared before the dump file is created. SSL session data is never written to disk.

#### § 7.8 SSL Inspection Process by Snooping

[0142] FIG. **15** is a flowchart of a process **680** for SSL (or other type of encrypted traffic) inspection by snooping, such as via a node operating as the snooping proxy **660**. The process **680** contemplates implementation as a method, as a computer-readable code stored on a non-transitory computer-readable storage medium for programming the node operating as the snooping proxy **660**, and the node operating as the snooping proxy **660**.

[0143] The process **680** includes monitoring traffic between a user device and the Internet (step **682**); detecting and monitoring a handshake between the user device and an endpoint for determining keys associated with encryption between the user device and the endpoint (step **684**); monitoring encrypted traffic between the user device and the endpoint subsequent to the handshake based on the keys (step **686**); and performing one or more security functions on the encrypted traffic based on the monitoring (step **688**). The node can be the enforcement node **150** that is part of a cloud-based security system, i.e., the cloud-based system **100**, and configured inline between the user device and the endpoint.

[0144] The process **680** can further include one of blocking or allowing the encrypted traffic based on the one or more security functions. The one or more security functions can include any of access control, threat prevention, and data protection, as described in detail herein. The endpoint can include an application utilizing certificate pinning. The process **680** can further include obtaining data related to the keys from a traffic-forwarding application executed on the user device. The process **680** can further include blocking the encrypted traffic responsive to being unable to decrypt the encrypted traffic with the keys.

#### § 8.0 Encrypted Tunnel Inspection with Customer CA

[0145] FIG. **16** is a diagram of encrypted tunnel inspection with the cloud-based system **100**. FIG. **16** includes a user device **300**, an enforcement node **150**, and an end server **404**, for illustrating MITM inspection. The user device **300** is configured to communicate with the end server **404**, with the enforcement node **150** being an inline device, e.g., proxy. The user device **300** sends either a CONNECT request (explicit proxy) or SSL request service name indication (SNI) (transparent proxy) to the enforcement node **150** (step **710-1**). The enforcement node **150** initiates SSL to the original URL (from the request) of the server **404** (step **710-2**). The server **404** responds to the enforcement node **150** with its certificate (step **710-3**). The enforcement node **150** generates and responds with a certificate (step **710-4**). Here, the enforcement node **150** has a trusted MITM certificate signed with the customer intermediate CA. This is where the cloud HSM fits—the enforcement node **150** makes a request to the cloud HSM to sign a MITM certificate's CSR (cert signing request) digest if it doesn't have the cert already cached in the data center.

[0146] The user device **300** validates the MITM certificate (step **710-5**), and SSL is negotiated between the user device **300** and the enforcement node **150** (step **710-6**). SSL is negotiated between the enforcement node **150** and the server **404** (step **710-7**). The enforcement node **150** now can view and inspect encrypted traffic between the user device **300** and the server **404**.

#### § 8.1 Key Safeguarding

[0147] FIG. **17** is a diagram of key storage **750** and key distribution **760**, in an embodiment, in the cloud-based system **100**. The key storage **750** can include a customer root certificate in a customer key store **570** that includes an HSM, a customer intermediate CA stored at the central authority **152** or in the cloud HSM, and a short-lived CA and a domain certificate (MITM) stored in memory at

the enforcement node **150**. Of note, the customer intermediate CA stored at the central authority **152** is stored on a disk. For the key distribution **760**, the central authority **152** sends the short-lived CA to the enforcement nodes **150** which use them to get domain certificates. In the cloud HSM model, the central authority **152** does not need to send any short-lived CA private to the enforcement nodes—that's what cloud HSM solves. Instead, the central authority sends a short lived authenticate token (oauth) to the cloud HSM service, which will, in turn, be used by the enforcement node to request the signing of a MITM certificate on the fly. FIG. **18** is a diagram illustrating certificates, their descriptions, and private key storage, based on FIG. **17**.

[0148] There can be various techniques to secure the intermediate CA private keys, but the crux of the problem is that the approach above stores the intermediate CA private key itself in a SQL DB and the short-lived issuing CA keys both on disk and in memory. Storing CA keys on disk or in memory is considered a risky design that FSI and banking customers cannot live with. Techniques for safeguarding can include, for example, auto rotating the short-lived CA to ensure the intermediate CA key exists for its entirety in a single server, only having short-lived keys on the enforcement nodes **150**, only having short-lived keys held in memory, minimal usage of the intermediate CA key at the central authority **152**, and the like. The technique can also include separation of duties—No CloudOps engineer access to code; No developer access to production (Central Authority); Stringent access to production systems—role-based access control with three factor authentication (ISO 27001, SOC-II); Logging and monitoring—extensive, of all commands executed on production host and jump servers (ISO 27001); File Integrity Monitoring—continuous deployment of critical files in case of unauthorized modification; and the like.

## § 8.2 Motivation and Requirements

[0149] 1) Every certificate private key issued from a customer's enterprise PKI and residing in the boundary of the cloud-based system **100**, must persist for its entire lifetime in a FIPS 140-2 Level 3 HSM. [0150] 2) The HSM has to have a hardware form factor And FIPS 140-2 Level 3 validated—not a “software HSM” [0151] 3) The feature has to provide a turnkey solution in which the cloud-based system **100** owns and manages the HSM. [0152] 4) Optionally, the solution can use a managed HSM solution by a public Cloud Vendor—not a dedicated HSM appliance that requires complex management and has high costs. [0153] 5) The design should be agnostic as possible to the underlying HSM vendor and form factor. It should be flexible enough to allow for possible integration with a customer managed HSM, through standard protocols such as REST, PKCS #11, KMIP. It should be flexible to “easily” swap between a different cloud HSM vendor in the future or even with an HSM appliance. [0154] 6) The solution should allow a customer to switch, without any downtime, between an existing software protection key and hardware based key protection.

## § 9.0 Cloud HSM System for MITM Inspection

[0155] FIG. **19** is a network diagram of the cloud-based system **100** and a cloud HSM system **800**. The cloud-based system **100** includes MITM domain certificates stored at the enforcement nodes **150**, and synchronized therebetween including in a memory cache **805**, as well as in different data centers. The cloud HSM system **800** can be a separate cloud system, i.e., a different cloud provider from the cloud-based system **100**, as well as part of the cloud-based system **100**. In this example, the cloud HSM system **800** is illustrated as a separate system, but those skilled in the art will recognize it can also be integrated. The cloud HSM system **800** includes a cloud key management system (KMS) **810** and HSM appliances **820** that are physical hardware.

[0156] FIG. **20** is a diagram illustrating workflow between the user device **300**, the enforcement node **150**, and the cloud HSM system **800**. Again, the design objective is that any certificate private key issued from the Enterprise PKI (Root or iCA) that has the power to issue subordinate certificates of any kind must be generated in and persist for its entire lifetime in a FIPS 140-2 Level 3 HSM, namely the HSM appliance **820** in the cloud HSM **800**. As such, the customer intermediate CA (CICA) is not stored at the central authority **152**, as in FIG. **17**. Also, the approach described here removes the need for the short-lived certificates.

[0157] The enforcement nodes **150** will make a request to the cloud HSM **800** only for certificates that are not in its cache. Once the certificate is created, it will be stored inside the cache for all subsequent requests to the same domain. A key assumption is that domain certificate keys (MITM) do NOT need to reside in the HSM.

[0158] FIG. **20** illustrates an example workflow which is segmented between the first access to a domain and subsequent access to the domain, e.g., example.com. For the first access to the domain, the user device **300** sends a hello to example.com which is intercepted by the enforcement node **150** (step **830-1**). The enforcement node **150** uses an RSA (Rivest-Shamir-Adleman) key pair and certificate signing request (CSR) (step **830-2**). The enforcement node **150** makes an asymmetric sign request to the cloud HSM **800** (step **830-3**), and the cloud HSM **800** returns a digital signature (step **830-4**). The enforcement node **150** merges the digital signature and certificate (step **830-5**) and completes the SSL handshake (step **830-6**). This certificate is then cached.

[0159] For subsequent access to the domain (step **830-7**), the enforcement node **150** fetches the cached certificate (step **830-8**) and completes the SSL handshake (step **830-9**).

[0160] Whenever the enforcement node **150** intercepts an SSL connection to a destination domain for which a MITM certificate does not already reside inside the local certificate/key cache, it will make a signing request for the MITM CSR to the HSM that holds the intermediate CA private key. The HSM, in turn, will respond with the digital signature for the certificate. In this methodology, the cleartext CA private key is used only within the FIPS boundaries of the HSM.

[0161] FIG. **21** is a diagram of interactions between the cloud-based system **100** and the cloud HSM **800**. Of note, there are two workflows **850**, **860** including CA enrollment **850** and MITM runtime **860**. The CA enrollment **850** is performed between the central authority **152**, a customer certificate authority **870**, and the cloud HSM **800**, including the cloud KMS **810** and the HSM appliance **820**. The CA enrollment **850** is performed up front, to enroll the customer key. The MITM runtime **860** is performed at runtime, and includes the example workflow of FIG. **20**. The cloud HSM include the cloud KMS **810** and the HSM appliance **820** which can communicate using PKCS #11, such as described in July 2020: OASIS PKCS #11 v3.0, the contents of which are incorporated by reference. The cloud-based system **100** can communicate to the cloud KMS **810** via a Representational state transfer (REST) API.

[0162] The CA enrollment **850** includes the central authority **152** requesting creation of a key from the cloud KMS **810** (step **850-1**). The cloud KMS **810** requests the HSM appliance **820** to create a key pair (step **850-2**) for the intermediate central authority. Also, the central authority **152** creates a customer certificate authority CSR (step **850-3**), sends a sign CSR-hash to the cloud HSM **810** (step **850-4**) which sends the sign CSR-hash to the HSM appliance for signing **820** (step **850-5**). The central authority **152** requests the customer certificate authority **870** to issue the certificate by signing the CSR (step **850-6**) and returns an X.509 certificate to the central authority **152** (step **850-6**).

[0163] The MITM runtime **860** is the same as the workflow of FIG. **20**. FIG. **22** is a diagram illustrating certificates, their descriptions, and private key storage, based on FIG. **21**.

[0164] FIG. **23** is a diagram of key storage **900** and key distribution **910** with the cloud HSM **800**, in an embodiment, in the cloud-based system **100**. Of note, it is likely that every data center for the cloud-based system **100** may not have a locally available HSM, so several regional hub Cloud HSMs **800-1**, **800-2** will provide service to the enforcement nodes **150** in the nearby, geographically, data centers. Since the throughput expectation for first seen domains is not high, especially in steady state, a regional hub model with central locations (US-West, US-East, US-Central, EU-West, EU-Central, EU-South, APAC) will suffice to ensure low latency <25 ms for the first seen domains and ensure minimal UX impact.

[0165] Inserting an HSM into the data path, for providing digital signing services, fundamentally changes the distributed cloud architecture design, in which each enforcement node **150** instance acts as an independent processing entity, to a design in which each enforcement node **150** needs to



make a round trip to the nearby HSM to get the signing services.

[0166] Rather than distributing the work as in the existing design, the main differences are that the enforcement node **150** will no longer act as an issuing CA, the short-lived CA layer is not needed, a new issuing CA service role will be introduced to make the enforcement node **150** agnostic of the preferred HSM form factor/model/interface, and provide a caching service.

#### § 9.1 Process

[0167] FIG. **24** is a flowchart of a process **1000** for man-in-the-middle inspection of encrypted traffic with cloud-based hardware security modules (HSM) with multi-tenant key infrastructure. The process **1000** contemplates implementation as a method, as a computer-readable code stored on a non-transitory computer-readable storage medium for programming the node **150**, etc.

[0168] The process **1000** includes, responsive (**1002**) to monitoring a user device (**300**), detecting a request for encrypted traffic to a domain from the user device (**300**); checking (**1004**) if a domain certificate for the domain is available in cache; responsive (**1006**) to the domain certificate being in the cache, creating a first tunnel (**610**) to the domain and a second tunnel (**612**) to the user device; and, responsive (**1008**) to the domain certificate not being in the cache, generating the domain certificate with a cloud hardware security module (HSM) system (**800**), and creating the first tunnel (**610**) and the second tunnel (**612**).

[0169] The process (**1000**) includes, responsive to generating the domain certificate, caching the domain certificate and synchronizing with other nodes (**150**) in the cloud-based system (**100**). The process (**1000**) includes monitoring the encrypted traffic for one or more security functions. The process (**1000**) includes one of blocking or allowing the encrypted traffic based on the one or more security functions. The process (**1000**) includes prior to the detecting, enrolling the domain in the cloud HSM (**800**) with a customer certificate authority (**870**). A cleartext private key associated with the customer certificate authority is constrained to an HSM.

[0170] The generating includes generating a key pair and certificate signing request; requesting the cloud HSM (**800**) asymmetrically sign; receiving a digital signature from the cloud HSM (**800**); and merging the digital signature and a certificate. The cloud HSM (**800**) can be a separate cloud system from the cloud-based system (**100**). The cloud HSM (**800**) includes one or more key management systems (KMSs) (**810**) connected to one or more physical HSMs (**820**). The first tunnel (**610**) and the second tunnel (**612**) are created based on a plurality of certificates including a customer root certificate that resides in a customer HSM, customer intermediate certificates previously stored in the cloud HSM, and the domain certificates issued at runtime or from the cache.

#### § 10.0 Conclusion

[0171] It will be appreciated that some embodiments described herein may include one or more generic or specialized processors (“one or more processors”) such as microprocessors; Central Processing Units (CPUs); Digital Signal Processors (DSPs); customized processors such as Network Processors (NPs) or Network Processing Units (NPUs), Graphics Processing Units (GPUs), or the like; Field Programmable Gate Arrays (FPGAs); and the like along with unique stored program instructions (including both software and firmware) for control thereof to implement, in conjunction with certain non-processor circuits, some, most, or all of the functions of the methods and/or systems described herein. Alternatively, some or all functions may be implemented by a state machine that has no stored program instructions, or in one or more Application-Specific Integrated Circuits (ASICs), in which each function or some combinations of certain of the functions are implemented as custom logic or circuitry. Of course, a combination of the aforementioned approaches may be used. For some of the embodiments described herein, a corresponding device in hardware and optionally with software, firmware, and a combination thereof can be referred to as “circuitry configured or adapted to,” “logic configured or adapted to,” etc. perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. on digital and/or analog signals as described herein for the various embodiments.

[0172] Moreover, some embodiments may include a non-transitory computer-readable storage medium having computer-readable code stored thereon for programming a computer, server, appliance, device, processor, circuit, etc. each of which may include a processor to perform functions as described and claimed herein. Examples of such computer-readable storage mediums include, but are not limited to, a hard disk, an optical storage device, a magnetic storage device, a Read-Only Memory (ROM), a Programmable Read-Only Memory (PROM), an Erasable Programmable Read-Only Memory (EPROM), an Electrically Erasable Programmable Read-Only Memory (EEPROM), Flash memory, and the like. When stored in the non-transitory computer-readable medium, software can include instructions executable by a processor or device (e.g., any type of programmable circuitry or logic) that, in response to such execution, cause a processor or the device to perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. as described herein for the various embodiments.

[0173] The foregoing sections include headers for various embodiments and those skilled in the art will appreciate these various embodiments may be used in combination with one another as well as individually. Although the present disclosure has been illustrated and described herein with reference to preferred embodiments and specific examples thereof, it will be readily apparent to those of ordinary skill in the art that other embodiments and examples may perform similar functions and/or achieve like results. All such equivalent embodiments and examples are within the spirit and scope of the present disclosure, are contemplated thereby, and are intended to be covered by the following claims. Moreover, it is noted that the various elements, operations, steps, methods, processes, algorithms, functions, techniques, etc. described herein can be used in any and all combinations with each other.

## Claims

1. A method for man-in-the-middle (MITM) inspection of encrypted traffic in a cloud-based security system, the method comprising: receiving, at a node within the cloud-based security system, a request from a user device to access encrypted traffic from a server; obtaining a domain certificate corresponding to the server; establishing a first encrypted tunnel between the node and the user device and a second encrypted tunnel between the node and the server using the domain certificate; and inspecting encrypted traffic communicated between the user device and the server.
2. The method of claim 1, wherein the domain certificate is generated using a cloud-based Hardware Security Module (HSM).
3. The method of claim 2, wherein the domain certificate is generated by: generating a certificate signing request (CSR) and a key pair; requesting asymmetric signing of the CSR from the cloud-based HSM; receiving a digital signature from the cloud-based HSM; and merging the digital signature with the CSR to form the domain certificate.
4. The method of claim 2, wherein a private key associated with an intermediate Certificate Authority (CA) used to sign the domain certificate is stored exclusively in the cloud-based HSM.
5. The method of claim 1, further comprising caching the domain certificate locally at the node for subsequent use.
6. The method of claim 5, further comprising synchronizing cached domain certificates between multiple nodes within the cloud-based security system.
7. The method of claim 1, wherein the inspecting encrypted traffic comprises performing security functions including threat prevention, data loss prevention (DLP), intrusion detection, malware detection, or web content filtering.
8. The method of claim 1, further comprising enrolling the node with a customer Certificate Authority (CA) through a cloud-based Hardware Security Module (HSM) prior to obtaining the domain certificate.
9. The method of claim 1, wherein the obtaining the domain certificate comprises retrieving the

domain certificate from a memory cache if previously generated and cached.

**10.** The method of claim 1, further comprising blocking or allowing encrypted traffic based on results of the inspecting.

**11.** The method of claim 1, wherein the establishing the first and second encrypted tunnels comprises negotiating Secure Sockets Layer (SSL) or Transport Layer Security (TLS) sessions.

**12.** The method of claim 11, wherein the inspecting encrypted traffic includes decrypting the traffic using session keys obtained through the SSL or TLS negotiation.

**13.** The method of claim 1, wherein the node operates as an inline proxy configured between the user device and the server.

**14.** The method of claim 1, wherein the node is configured to transparently intercept the request from the user device without explicit proxy configuration at the user device.

**15.** The method of claim 1, wherein the cloud-based security system comprises multiple geographically distributed nodes, each configured to perform MITM inspection.

**16.** The method of claim 1, further comprising: detecting failure conditions where encrypted traffic cannot be decrypted; and responsive to the detecting, blocking the encrypted traffic.

**17.** The method of claim 1, wherein the domain certificate is an intermediate certificate signed by a customer-specific root Certificate Authority (CA).

**18.** The method of claim 1, wherein the request from the user device includes Server Name Indication (SNI) information identifying the server.

**19.** The method of claim 1, wherein inspecting encrypted traffic further comprises enforcing granular policies based on user identity, URL category, or application type.

**20.** A node in a cloud-based security system configured for man-in-the-middle (MITM) inspection of encrypted traffic, the node comprising: one or more processors and memory storing instructions that, when executed, cause the one or more processors to: receive a request from a user device to access encrypted traffic from a server; obtain a domain certificate corresponding to the server; establish a first encrypted tunnel between the node and the user device and a second encrypted tunnel between the node and the server using the domain certificate; and inspect encrypted traffic communicated between the user device and the server.

---