

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250260630

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

BEHL; Sunny et al.

PLATFORM FOR AUTOMATED MANAGEMENT OF SERVERS

Abstract

Presented herein are systems and methods for managing networked environments. A computer system may provide a user interface for configuring a plurality of groups of servers each hosting a resource for an application. The user interface may include: a first element configured to select, upon interaction, at least one group of servers from the plurality of groups of servers to which to install a patch for the application; a second element configured to identify, upon interaction, a first group of servers of the plurality of groups of servers to which to transfer network traffic associated with the application and communicated with a second group of servers of the plurality of groups of servers; and a third element configured to provide, upon interaction, one or more performance indicators for at least one of a plurality of functions of the application.

Inventors: BEHL; Sunny (Irving, TX), BILLA; Suresh Kanyaka (Frisco, TX), RAVI; Sandeep (Coppell, TX)

Applicant: Citigroup Technology, Inc. (New York, NY)

Family ID: 86993276

Assignee: Citigroup Technology, Inc. (New York, NY)

Appl. No.: 19/191439

Filed: April 28, 2025

Related U.S. Application Data

parent US continuation 18394929 20231222 parent-grant-document US 12294509 child US 19191439

parent US continuation 18215598 20230628 parent-grant-document US 11863308 child US 18394929

parent US continuation 18099826 20230120 parent-grant-document US 11695655 child US 18215598

Publication Classification

Int. Cl.: H04L41/50 (20220101); G06F11/32 (20060101); G06F11/34 (20060101)

U.S. Cl.:

CPC H04L41/5096 (20130101); G06F11/323 (20130101); G06F11/3409 (20130101);

Background/Summary

CROSS REFERENCES TO RELATED APPLICATIONS [0001] The present application claims priority under 35 U.S.C. § 120 as a continuation of U.S. patent application Ser. No. 18/394,929, titled “Platform for Automated Management of Servers,” filed Dec. 22, 2023, which claims priority under 35 U.S.C. § 120 as a continuation of U.S. patent application Ser. No. 18/215,598, titled “Platform for Automated Management of Servers,” filed Jun. 28, 2023, now U.S. Pat. No. 11,863,308 which claims priority under 35 U.S.C. § 120 as a continuation of U.S. patent application Ser. No. 18/099,826, titled “Platform for Automated Management of Servers,” filed Jan. 20, 2023, now U.S. Pat. No. 11,695,655, each of which are incorporated by reference in their entirety.

TECHNICAL FIELD

[0002] This application generally relates to managing servers in networked environments. In particular, the present application relates to a graphical user interface for managing server groups in networked environments.

BACKGROUND

[0003] Servers physically residing across various sites may host and provide resources for an application (e.g., a web application) over a computer networked environment. The application may have a plethora of functionality to be accessed by end-user devices. In doing so, the application may produce an immense amount of complex data related to the various operations of the application, such as network metrics, application performance, and server. Due to the enormous quantity of data, it may be difficult for a network administrator to manage the server's hosted resources for the application as well as the application itself. Furthermore, the network administrator may have to access one tool to obtain one set of data and another tool to retrieve another set of data to make any decisions regarding the system. Due to the inability to manage the network, any problems affecting the performance of the application or the servers hosting the application may remain unresolved and continue to linger.

SUMMARY

[0004] Disclosed herein are systems and methods for managing networked environments. Servers residing at different sites and hosting resources for an application may collect an immense amount of data associated with the application. The data may include, for example, the network traffic from the server to the end-user devices, invocations of various functions of the application, and information about the end-user devices accessing the application, among others. The amount of data collected by the servers may become enormous and complex, when aggregated over multiple servers, applications, and end-user devices over the networked environment.

[0005] A network administrator of the environment may use the data to diagnose any problems within the servers or the application and carry out any measures in an attempt to address the issues or sub-optimal conditions. The administrator may rely on a multiplicity of tools (e.g., applications or services) to view pieces of the data to assess the environment and to undertake any actions based on such assessments. For example, to update a version of the application hosted on a given group

of servers, the administrator may access the individual servers to find which version of the application is installed thereon. Upon finding the version, the administrator may have to run a shutdown process for the individual servers and open a tool to roll out and install the update to the application. Once the installation is complete, the administrator may have to run yet another tool to perform validation check of the update. Other aspects of managing the environment may similarly rely on multiple tools, such as for failover management, predictive analytics, reliability checks, and incident ticketing, among others.

[0006] With more and more tools for various, specific aspects of the networked environment, it may be overly cumbersome for the network administrator to effectuate effective administration and management of the servers in the networked environment. The inability to quickly access and digest data relevant to management of the network may leave problems and other sub-optimal conditions in the environment unaddressed. This may result in wasted consumption of computing resources at each server, reduced capabilities of the applications hosted thereon, and degradation of network conditions, such as lower available bandwidth and increased latency. The culmination of these sub-optimal conditions may lead to the deterioration of the overall performance of the environment and a decrease in the quality of human-computer interaction (HCI) between the end-users and the applications hosted on the servers.

[0007] To address these and other technical challenges, an automated management service may provide a user interface (also referred herein as a dashboard) to integrate various aspects of network administration. The user interface may provide a set of options to retrieve performance data of applications and servers in the environment as well as to carry out various actions to manage the applications and servers in the network. The options provided by the user interface may include, for example, a one-touch patch management, a one-touch continuity of traffic fail over, predictive analytics, site reliability, and a services assistant, among others.

[0008] The one-touch patch management may provide for automated process of patch installation for an application hosted on servers with a single interaction with a user interface element of the dashboard. This feature may eliminate manual steps and system administration resources doing repeatable work from patch scheduling phase to patch execution via orchestration. The patch management may automate shutting down of servers before triggering patch and restoring same set of servers after patching, and may provide for the ability to reboot servers and perform post patching checks including validation, without any manual intervention.

[0009] The one-touch continuity traffic failover may provide self-service mechanisms to failover application traffic between groups of servers at different sites (e.g., data centers), during maintenance windows and during major incidents. This may mitigate the impacts of any interruption in server to end-users and may allow quick transfer of traffic from one site to another site until restoration of the original group of servers. The failover feature may enable integration with server infrastructure to automate datacenter failover and may provide for live application traffic trends across sites.

[0010] The predictive analytics interface may provide for observability of applications and associated functionality (sometimes herein referred to as transactions) by displaying real-time performance of services. Machine learning models may be used to quickly process accumulated data to produce a performance indicator for each function of a given application hosted on a server. This solution may involve automation leveraging historical application health metrics and anomaly detection techniques comparing dynamic baselines with observed values. The analytics interface may also provide a drill down of historical trends of various related metrics, such as volumes, slow calls, and success or error rates, among others.

[0011] In addition, the user interface may provide for a site reliability interface. This feature allows the user to define a service level indicator (SLI) across various specifications, such as availability, latency, quality, coverage, and durability, among others. The automated management service may record success and failure counts for each function (or transaction type or feature) provided by the

application. The service may calculate availability and error based on a defined service level objection (SLO). This site reliability interface may also provide ability to select different on-boarded applications along with types of functions to present reliability metrics in a selected timeframe. For each type of function, the site reliability interface may provide real-time incremental SLO and error rates indicating upwards or downwards trend. The feature may provide comprehensive view of all types of functions, with rates of failure, availability, and error budgets, among others.

[0012] The user interface may also provide an intelligent customer services assistant. The self-service feature may include a services agent to perform application health checks and provide information from a knowledge base around frequently reported issues. The agent may also provide for a ticketing service to raise incidents and contacting service agents to allow for a seamless experience in opening tickets and post initial validations. This feature may also prevent duplicate incidents to be raised by multiple agents for the same issue, and may enable incidents to be opened with appropriate quality with minimal information auto filled via questionnaire interface.

[0013] In this manner, the management service may provide for automation of tactical activities with strategic oversight, instead of manually spending time and skilled expertise on issues within the network. The dashboard user interface may thus augment resource capabilities, allowing the network administrator to focus on issues. The automated management service may provide for reduction of operating maintenance by improving efficiency and significantly reducing the manual and tactical resource intensive efforts with automated workflows.

[0014] The service may improve mean time to recovery (MTTR) by eliminating hosting batches and team involvement in tedious issues when log analysis and re-runs may be used to address such issues. The service may reduce noise alerts by filtering alerts and suppressing false positives. Furthermore, the automated management service may provide for process simplification and workflow automation by improving visibility to various processes, indication of service level agreement (SLA) objectives, and certificate and risk management. The improvement may involve reduced risk in missing SLA objectives with heat maps and a forecast of potential SLA deviations.

[0015] The service may also allow prediction, prescription, and prevention of outages using combination of artificial intelligence (AI) and automation, thereby ensuring predictable and silent operations. The service may minimize impacts due to change in implementations through environment analysis, automated severity assessments, self-diagnosis, and automated health check and validations, among others. The service may provide for proactive optimizations through one-touch failover, identification of repetitive issues, and anomaly detection, thereby providing an opportunity to pinpoint root causes and proactively prescribe countermeasures. The service may construct end-to-end mapping of impact to function types, measuring latency, slow calls, availability, and errors, among others. The automated management service may transform operations to support autonomous operations by eliminating tickets through the services agent and knowledge database.

[0016] Aspects of present disclosure are directed to systems, methods, and non-transitory computer readable media for managing networked environments. A computer system having one or more processors coupled with memory may provide a user interface for configuring a plurality of groups of servers each hosting a resource for an application. The user interface may include: a first element configured to select, upon interaction, at least one group of servers from the plurality of groups of servers to which to install a patch for the application; a second element configured to identify, upon interaction, a first group of servers of the plurality of groups of servers to which to transfer network traffic associated with the application and communicated with a second group of servers of the plurality of groups of servers; and a third element configured to provide, upon interaction, one or more performance indicators for at least one of a plurality of functions of the application.

[0017] The computing system may execute, responsive to a first interaction with the first element, a

patch management process by shutting down the selected at least one group of servers, installing the patch for the application on the at least one group of servers, and validating installation of the patch. The computing system may perform, responsive to a second interaction with the second element, a traffic fail over process by transferring the network traffic associated with the application from the second group of servers to the first group of servers. The computing system may provide, responsive to a third interaction with the third element, a performance indicator for a function of the plurality of functions of the application based on rolling historical data for the function.

[0018] In one embodiment, the computing system may apply the rolling historical data of the function to a machine learning (ML) model to determine the performance indicator identifying a predicted likelihood of the application successfully carrying out the function. In another embodiment, the computing system may select, from a plurality of performance indicators, the performance indicator based on a comparison between a performance metric for the function and a threshold value. In yet another embodiment, the computing system may determine the threshold value with which to compare against the performance metric for the function based on the historical rolling data for the function.

[0019] In yet another embodiment, the user interface may include a fourth element configured to provide, upon interaction, metrics over time for the function selected from the plurality of functions. The computing system may generate a plurality of performance metrics over time based on the historical rolling data used to provide the performance indicator.

[0020] In yet another embodiment, the user interface may include a fourth element configured to provide, for each of the plurality of groups of servers, a respective status indicator identifying whether the patch has been installed on a corresponding group of servers. In yet another embodiment, the user interface may include a fourth element configured to provide, for each of the plurality of groups of servers, a respective statistic indicator on the network traffic associated with the application at a corresponding group of servers.

[0021] In yet another embodiment, the user interface may include a fourth element configured to set, upon interaction, a schedule for executing the patch management process of the at least one group of servers. The computing system may execute the patch management process in accordance with the schedule set using the fourth element. In yet another embodiment, the computing system may identify a first stack in the first group of servers corresponding to a second stack in the second group of servers configured to host the application. In yet another embodiment, the plurality of groups of servers may be arranged by at least one of: application type, geographic location, or serving device type.

[0022] It is to be understood that both the foregoing general description and the following detailed description are exemplary and explanatory and are intended to provide further explanation of the embodiments described herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0023] The accompanying drawings constitute a part of this specification, illustrate an embodiment, and together with the specification, explain the subject matter of the disclosure.

[0024] FIG. 1 depicts a block diagram of a system for managing networked environments in accordance with an illustrative embodiment;

[0025] FIG. 2 depicts a screenshot of an dashboard user interface of the system for managing networked environments in accordance with an illustrative embodiment;

[0026] FIG. 3 depicts a block diagram of a system for patch management in networked environments in accordance with an illustrative embodiment;

[0027] FIG. 4 depicts a screenshot of an execution interface in the dashboard user interface for the system for patch management in accordance with an illustrative embodiment;

[0028] FIG. 5 depicts a screenshot of a status interface in the dashboard user interface for the system for patch management in accordance with an illustrative embodiment;

[0029] FIG. 6 depicts a block diagram of a system for failover management in networked environments in accordance with an illustrative embodiment;

[0030] FIG. 7 depicts a screenshot of a network health interface in the dashboard user interface for the system for failover management in accordance with an illustrative embodiment;

[0031] FIGS. 8A and 8B depict screenshots of an execution interface in the dashboard user interface for the system for failover management in accordance with an illustrative embodiment;

[0032] FIG. 9 depicts a screenshot of a traffic pattern interface in the dashboard user interface for the system for failover management in accordance with an illustrative embodiment;

[0033] FIG. 10 depicts a block diagram of a system for performance analytics in networked environments in accordance with an illustrative embodiment;

[0034] FIG. 11 depicts a screenshot of a performance indicator interface in the dashboard user interface for the system for performance analytics in accordance with an illustrative embodiment;

[0035] FIG. 12 depicts a screenshot of a drill down interface in the dashboard user interface for the system for performance analytics in accordance with an illustrative embodiment;

[0036] FIG. 13 depicts a block diagram of a system for site reliability evaluation in networked environments in accordance with an illustrative embodiment;

[0037] FIG. 14 depicts a screenshot of an on-boarding interface in the dashboard user interface for the system for site reliability evaluation in accordance with an illustrative embodiment;

[0038] FIG. 15 depicts a screenshot of a reliability indication interface in the dashboard user interface for the system for site reliability evaluation in accordance with an illustrative embodiment;

[0039] FIG. 16 depicts a screenshot of a drill down interface in the dashboard user interface for the system for site reliability evaluation in accordance with an illustrative embodiment;

[0040] FIG. 17 depicts a screenshot of a heat map interface in the dashboard user interface for the system for site reliability evaluation in accordance with an illustrative embodiment;

[0041] FIG. 18 depicts a block diagram of a system for services assistance in a system for managing networked environments in accordance with an illustrative embodiment;

[0042] FIG. 19 depicts a screenshot of a health check interface in the dashboard user interface for the system for services assistance in accordance with an illustrative embodiment;

[0043] FIG. 20 depicts a screenshot of a query interface in the dashboard user interface for the system for services assistance in accordance with an illustrative embodiment; and

[0044] FIG. 21 depicts a flow diagram of a method of managing networked environments in accordance with an illustrative embodiment.

DETAILED DESCRIPTION

[0045] Reference will now be made to the embodiments illustrated in the drawings, and specific language will be used here to describe the same. It will nevertheless be understood that no limitation of the scope of the disclosure is thereby intended. Alterations and further modifications of the features illustrated here, and additional applications of the principles as illustrated here, which would occur to a person skilled in the relevant art and having possession of this disclosure, are to be considered within the scope of the disclosure.

[0046] The present disclosure is directed to systems and methods for managing networked environments. An automated management service may aggregate data from servers hosting applications accessed by end-user devices and provide a dashboard user interface to administer and manage various operations of the network environment. The dashboard user interface may include: a patch installation element to initiate a patch management process to at least one set of servers; a failover execution element to carry out a network traffic failover from one set of servers to another

set of servers; and a predictive analytics element to provide a set of performance indicators for various functions of a given application, among others. In this manner, the automated management service may allow the network administrator to quickly retrieve desired data (e.g., statistics and performance metrics) about the network and promptly take proper actions to manage various aspects of network operations.

[0047] FIG. 1 depicts a block diagram of a system **100** for managing networked environments. In overview, the system **100** may include at least one automated management service **102**, one or more servers **104A-1** to **104N-X** (hereinafter generally referred to as servers **104**) arranged, situated, or distributed across a set of server groups **106A-N** (hereinafter generally referred to as server groups **106**); and at least one database **108**, among others, communicatively coupled with one another via at least one network **110**. The automated management service **102** may include at least one dashboard handler **112**, at least one patch manager **114**, at least one failover manager **116**, at least one analytics evaluator **118**, at least one reliability evaluator **120**, and at least one assistant handler **122**, among others. The automated management service **102** may provide at least one user interface **124**. The user interface **124** may include at least one patch installer user element (UI) element **126**, at least one failover execution UI element **128**, at least one analytics retrieval UI element **130**, at least one reliability retrieval UI element **132**, and at least one assistant invocation UI element **134**, among others. The one or more servers **104** in at least one server group **106** may include or host at least one application **136**.

[0048] Various hardware and software components of one or more public or private networks **110** may interconnect the various components of the system **100**. Non-limiting examples of such networks may include Local Area Network (LAN), Wireless Local Area Network (WLAN), Metropolitan Area Network (MAN), Wide Area Network (WAN), and the Internet. The communication over the network may be performed in accordance with various communication protocols, such as Transmission Control Protocol and Internet Protocol (TCP/IP), User Datagram Protocol (UDP), and IEEE communication protocols, among others.

[0049] Each server **104** may be any computing device comprising one or more processors coupled with memory and software and capable of performing the various processes and tasks described herein. Each server **104** may be in communication with one another, one or more end-user customer devices, the automated management service **102**, and the database **108**, among others via the network **110**. The server **104** may be situated, located, or otherwise associated with at least one server group **106**. Each server group **106** may correspond to a data center, a branch office, or a site at which a subset of servers **104** is situated or associated. For instance, the first server group **106A** may correspond to a data center at first site including a first set of servers **104A-1** to **104A-X** and the second server group **106B** may correspond to a branch office at a second side including a second set of servers **106B-1** to **106B-X**.

[0050] At least one of the servers **104** may maintain, include, or otherwise host resources for the application **136**. The application **136** may be a cloud-based application (e.g., a Software as a Service (SaaS)) or a web application, among others, accessed by end-user customer devices that are communicatively coupled with the network **110**. For example, the application **136** may be an online banking application, a word processor, a spreadsheet program, a multimedia player, a video game, or a software development kit, among others. In some embodiments, the server **104** may include or run the application **136** itself. For instance, the server **104** may maintain or run virtual machine to run instances of the application **136** to be accessed by the end-user customer devices. In some embodiments, the servers **104** may be grouped, associated with one another, or otherwise arranged into the server group **106** by: a type of the application **136** provided; a geographic location; a network location; or a type of device (e.g., the end-user customer devices such as mobile phones, laptops, or desktops), among others.

[0051] The database **108** may store and maintain various data associated with the servers **104** across the server groups **106** and with the application **136** hosted thereon, among others. The

database **108** may also include a database management system (DBMS) to arrange and organize the data maintained thereon. The data may be produced from the application **136** running on the servers **104** and accessed by end-user customer devices over the network **110**. Upon production, the servers **104** (or the end-user customer devices) may store the data onto the database **108**. For instance, the database **108** may store and maintain a transaction log identifying communications exchanged over the network **110**, such as between end-user customer device and the servers **104**. The database **108** may store and maintain a process log for a given application **136** identifying functions, events, or otherwise component invoked in the application **136** when running on the servers **104** or when accessed by the end-user customer devices. The data maintained on the database **108** may be accessed by the automated management service **102**.

[0052] The automated management service **102** may be any computing device comprising one or more processors coupled with memory and software and capable of performing the various processes and tasks described herein. The automated management service **102** may be in communication with the servers **104** across different server groups **106**, one or more end-user customer devices, and the database **108**, among others via the network **110**. Although shown as a single component, the automated management service **102** may include any number of computing devices. For instance, the dashboard handler **112**, the patch manager **114**, the failover manager **116**, the analytics evaluator **118**, the reliability evaluator **120**, and the assistant handler **122** may be executed across one or more computing systems.

[0053] Within the automated management service **102**, the dashboard handler **112** may provide the user interface **124** for display of a computing device communicatively coupled with the automated management service **102**. The patch manager **114** may execute a patch installation process for the application **136** hosted on the servers **104**. The failover manager **116** may perform a traffic failover process from one server group **106** to another server group **106**. The analytics evaluator **118** may provide performance indicators for various functions of a given application. The reliability evaluator **120** may provide reliability indicators in accordance with service level objectives (SLOs). The assistant handler **122** may provide a service assistant to handle health checks and incident ticking for the application **136** or any of the servers **104**.

[0054] The user interface **124** may be a graphical user interface (GUI), with one or more elements to invoke various functions of the automated management service **102**. Upon interaction, the patch installer UI element **126** may invoke the patch manager **114** to execute the patch installation process. The failover execution UI element **128** may invoke the failover manager **116** to perform a traffic failover process from one server group **106** to another server group **106**. The analytics retrieval UI element **130** may invoke the analytics evaluator **118** to provide performance indicators for various functions of the application **136**. The reliability retrieval UI element **132** may invoke the reliability evaluator **120** to provide reliability indicators. The assistant invocation UI element **134** may invoke the assistant handler **122** to provide a service assistant to handle health checks and incident ticking. The user interface **124** may be for example in the manner depicted in FIG. 2.

[0055] FIG. 2 depicts a screenshot of a dashboard user interface **200** of the system for managing networked environments. In the depicted example, the dashboard user interface **200** may include a set of elements generally in the middle of the interface, such as: a first UI element **202** to open incident management (e.g., including failover management); a second UI element **204** to open site reliability measurements; a third UI element **206** to access predictive analytics; a fourth UI element **208** to access patch management; and a fifth UI element **210** to invoke the intelligent customer service agent, among others. The dashboard user interface **200** may also include at least one sixth element **212** to provide notifications and updates. The dashboard user interface **200** may include other UI elements, such as a tool bar along the top and access to other dashboards along the left.

[0056] FIG. 3 depicts a block diagram of a system **300** for patch management in the networked environment. The system **300** may include at least one automated management service **302**, one or more servers **304A-X** (hereinafter generally referred to as servers **304**) in a server group **306**, and at

least one database **308**, among others, communicatively coupled with one another via at least one network **310**. The automated management service **302** may include at least one dashboard handler **312** and at least one patch manager **314**, among others, and may provide at least one user interface **324**. The user interface **324** may include at least one patch installer UI element **326** and the patch management UI element **328**, among others. At least one of the servers **304** may host resources for at least one application **336**. The automated management service **302** may be used to facilitate installation of at least one patch **338** to the application **336**. Embodiments may comprise additional or alternative components or omit certain components from those of FIG. 3, and still fall within the scope of this disclosure.

[0057] The dashboard handler **312** executing on the automated management service **302** may provide the user interface **324** including the patch installer UI element **326**. The user interface **324** may be rendered, displayed, or otherwise presented on at least one display of the automated management service **302** or communicatively coupled with the automated management service **302**. The user interface **324** may be in the manner, for example, depicted in FIGS. 4 and 5, among others. The patch installer UI element **326** may correspond to at least one graphical user interface (GUI) element within the user interface **324**, such as a command button, a slider, a toggle switch, an image, a window, a prompt, or a container, among others, or any combination thereof. The dashboard handler **312** may monitor for an interaction with the patch installer UI element **326**, such as a mouse click, a screen touch, a key press, voice command, or a corresponding gesture, among others. The patch installer UI element **326** may identify or select at least one server group **306** from a set of groups of servers on which to set up install at least one patch **338** for the application **336** hosted at least one of the servers **304** in the server group **306**. In response to detecting an interaction with the patch installer UI element **326**, the dashboard handler **312** may call, invoke, or otherwise execute the patch manager **314**.

[0058] In some embodiments, the dashboard handler **312** may provide the user interface **324** to include at least one patch management UI element **340**. The patch management UI element **340** may correspond to a set of graphical user interface (GUI) elements within the user interface **324**, such as a radio button, a check box, a toggle switch, a text box, an image, a window, a prompt, or a container, among others, or any combination thereof. In some embodiments, the patch management UI element **340** may be part of the patch installer UI element **326** within the user interface **324**. For example, the dashboard handler **312** may present the patch management UI element **340** corresponding to a subset of constituent GUI elements of the patch installer UI element **326**. In some embodiments, the patch management UI element **340** may be separate from the patch installer UI element **326** within the user interface **324**. For instance, the dashboard handler **312** may present the patch management UI element **340** on a window or webpage separate from the patch installer UI element **326**.

[0059] In some embodiments, the patch management UI element **340** may include or provide information associated with the application **336**, the servers **304**, and server groups **306**, among others. The patch management UI element **340** may provide information relevant to the installation of the patch **338** for the application **336**. For instance, the patch management UI element **340** may include: a version identifier for the application **336** currently installed on each server **304** or server group **306**; an identifier for each server **304** or server group **306** on which the application **336** is installed; a status indicating a progress (e.g., downloading, setting up, validation, and completion) of installation of the patch **338** at the respective server **304** or server group **306**; and a time stamp for the status, among others.

[0060] In some embodiments, the patch management UI element **340** may accept, gather, or otherwise receive one or more parameters for the installation of the patch **338** for the application **336**. The parameters may define a set of stages of the setup process of the patch **338** on the server **304** or the server group **306** for the application **336** and may be entered by a user (e.g., a network administrator) of the automated management service **302**. The parameters may include, for

example: an application identifier referencing the application **336**; a version identifier corresponding to the patch **338** for the application **336** to be installed; an identification of a selected server **304** or server group **306** on which to carry out installation of the patch **338**; and a schedule for the installation of the patch **338** for the application **336**, among others. The schedule may define a time at which to carry out the stage of the installation process of the patch **338**. In invoking the patch manager **314**, the dashboard handler **312** may pass the parameters inputted into the patch management UI element **340**.

[0061] The patch manager **314** executing on the automated management service **302** may carry out, perform, or otherwise execute a patch management process, in response to the interaction on the patch installer UI element **326**. The patch management process may include the set of stages of installing the patch **338** for the application **336** whose resources are hosted on the selected server **304** or server group **306**. The patch management process may start from a shutting down of the servers **304** hosting resources for the application **336** in the selected server group **306**, setting up or installing the patch **338** for the application **336** on the servers **304**, and validating the installation of the patch **338**, among others. The patch manager **314** may carry out the patch management process in accordance with the defined parameters upon invocation from interaction with the patch installer UI element **326**, with minimal or no subsequent user interaction. In this manner, the patch manager **314** may automate the various stages of installation of the patch **338** to reduce manual human involvement.

[0062] In carrying out the patch management process, the patch manager **314** may retrieve, obtain, or otherwise identify the patch **338** to be installed. The patch **338** may be stored and maintained in a storage (e.g., the database **308** as depicted) accessible to the automated management service **302**. The patch **338** may define, identify, or otherwise include a set of updates to be applied to the application **336**. The updates included in the patch **338** may include, for example, addition of new functions, removal of previously provided functions, or modifications to existing functions in the application **336**. In some embodiments, the patch manager **314** may identify which patch **338** is to be installed, using the inputs from the patch installer UI element **326** and the patch management UI element **340** of the user interface **324**.

[0063] With the identification, the patch manager **314** may send, transmit, or otherwise provide the patch **338** to the selected server **304** or server group **306**. In conjunction, the patch manager **314** may run a shutdown sub-process the selected server **304** or server group **306** hosting resources for the application **336**. The shutdown sub-process may entail, involve, or include causing the servers **304** in the server group **306** offline to cease further access by the end-user customer devices. Upon shutting down, the patch manager **314** may carry out setting up or installation of the patch **338** for the application **336**. To set up, the patch manager **314** may run or execute the patch **338** to apply the set of updates to the application **336**, for example, by changing executable binary files corresponding to the application **336**.

[0064] Continuing on, the patch manager **314** may perform a validation sub-process (sometimes herein referred to as a post-patch check) on the installation of the patch **338**. In performing, the patch manager **314** may determine whether the patch **338** is successfully installed on the server **304** or server group **306**, without affecting other processes on the server **304** or server group **306**. When the installation is unsuccessful, the patch manager **314** may return an indication for presentation on the user interface **324** (e.g., via the patch management UI element **340**). In some embodiments, the patch manager **314** may perform the stages of the patch management process (e.g., the shutdown, setup, and validation) in accordance with the defined parameters. For example, the patch manager **314** may carry out individual sub-processes in accordance with the times identified by the schedule defined using the patch management UI element **340**.

[0065] FIG. 4 depicts a screenshot of an execution interface **400** in the dashboard user interface for the system for patch management. In the depicted example, the execution interface **400** may include a list of applications **402** to indicate statuses of patch installations. The list **402** may include

various information about the patch status, such as: a product name **404** to identify a type of application; a data center name **406** to identify a server group hosting the application; an operating system type **408** to identify the operating system at the server group; a validation status **410** identifying progress or completion of validation of the patch installation; and a patch status **412** identifying a progress of the overall patch installation process for the given application. The execution interface **400** may be used by the network administrator to view the patch status of various application across multiple sites.

[0066] FIG. 5 depicts a screenshot of a status interface **500** in the dashboard user interface for the system for patch management. In the depicted example, the status interface **500** may include various information on the status of patch installation management across multiple server groups or sites. The status interface **500** may include: a patch count interface **502** identifying a number of patches installed by sites or server groups; a patch velocity interface **504** identifying a rate at which the patch management process is successfully carried out by sites or server groups; and a patch status interface **506** identifying a number of patch management processes that were either successful or failed. The status interface **500** may be used by the network administrator to view statistics regarding the patch installation process across multiple sites for a given application.

[0067] FIG. 6 depicts a block diagram of a system **600** for failover management in networked environments. The system **600** may include at least one automated management service **602**, one or more servers **604A-1** to **604B-X** (hereinafter generally referred to as servers **604**) across at least two server groups **606A** and **606B** (hereinafter generally referred to as server group **606**), and at least one database **608**, among others, communicatively coupled with one another via at least one network **610**. The automated management service **602** may include at least one dashboard handler **612** and at least one failover manager **614**, among others, and may provide at least one user interface **624**. The user interface **624** may include at least one failover execution UI element **628** and at least one network statistics UI element **640**, among others. At least one of the servers **604** in each server group **606** may host respective resources for at least one application **636**. The server group **606** may facilitate network traffic **638** for communications between the end-user customer devices and the servers **604** to access resources for the application **636**. Embodiments may comprise additional or alternative components or omit certain components from those of FIG. 6, and still fall within the scope of this disclosure.

[0068] The dashboard handler **612** executing on the automated management service **602** may provide the user interface **624** including the failover execution UI element **628**. The user interface **624** may be rendered, displayed, or otherwise presented on at least one display of the automated management service **602** or communicatively coupled with the automated management service **602**. The user interface **624** may be in the manner, for example, depicted in FIGS. 7-9, among others. The failover execution UI element **628** may correspond to at least one graphical user interface (GUI) element within the user interface **624**, such as a command button, a slider, a toggle switch, a radio button, a check box, a text box, an image, a window, a prompt, or a container, among others, or any combination thereof.

[0069] The failover execution UI element **628** may be used to identify or select one server group **606B** to which to transfer network traffic **640** associated with the application **636** from another server group **606A**. In some embodiments, the failover execution UI element **628** may accept, gather, or otherwise receive one or more parameters for transferal of the network traffic **640** as part of the failover. The parameters may identify or include, for example, an application identifier corresponding to the application **636**; a source identifier corresponding to the server group **606** from which the network traffic **640** is to be transferred (e.g., the server group **606A** as depicted); and a destination identifier corresponding to the server group **606** to which the network traffic **640** is to be transferred (e.g., the server group **606B** as depicted), among others.

[0070] The dashboard handler **612** may monitor for an interaction with the failover execution UI element **628** to initiate the failover, such as a mouse click, a screen touch, a key press, voice

command, or a corresponding gesture, among others. The interaction may indicate a command to initiate the failover. In some embodiments, the dashboard handler **612** may also handle one or more interactions to the failover execute UI element **634** to enter or input the parameters defining the traffic failover. In response to detecting the interaction with the failover execution UI element **628** to initiate the failover, the dashboard handler **612** may call, invoke, or otherwise execute the failover manager **614**. In invoking, the dashboard handler **612** may pass the input parameters to the failover manager **614**.

[0071] The failover manager **616** executing on the automated management service **602** may carry out, perform, or otherwise execute a traffic failover process, in response to the interaction with the failover execution UI element **634**. The traffic failover process may correspond to or include moving, switching, or otherwise transferring the network traffic **638** from one server group **606** to another server group **606** (e.g., from the first server group **606A** to the second server group **606B** as depicted). The network traffic **638** may have been previously communicated with the server group **606** (e.g., the first server group **606A**) in providing end-user customer devices access to resources for the application **636**. The failover manager **616** may execute the traffic failover process in accordance with the parameters input via the failover execution UI element **634**.

[0072] From the parameters, the failover manager **616** may select or identify the server group **606** (e.g., the first server group **606A**) referenced by the source identifier from which the network traffic **638** is to be transferred. In some embodiments, the failover manager **616** may find, select, or otherwise identify at least one stack (e.g., a subset of servers **606**) within the server group **606** hosting resources for the application **636**. The failover manager **616** may find, select, or otherwise identify the network traffic **638** associated with the application **636** hosted on one or more servers **604** of the identified server group **606** (e.g., the first server group **606A**). The network traffic **638** may identify or include communications (e.g., data packets) exchanged between the servers **604** of the server group **606** and the end-user consumer devices in accessing the application **636**.

[0073] In addition, the failover manager **616** may identify the server group **606** (e.g., the second server group **606B**) to which the network traffic **638** is to be transferred. In some embodiments, the failover manager **616** may find, select, or otherwise identify at least one stack (e.g., a subset of servers **604**) within the server group **606** corresponding to the stack in the other server group **606** to which the network traffic **638** is to be transferred. The stack may correspond to a subset of servers **604** already hosting resources for the application **636** or another instance of the application **636** hosted thereon. The stack may also correspond to the subset of servers **604** in the server group **606** with availability to handle such resources **636** or communications with the end-user consumer devices to access the application **636**.

[0074] With the identifications, the failover manager **616** may instruct, command, or otherwise cause the servers **604** of the server group **606** (e.g., the server group **606A**) to move or transfer the network traffic **638** to the servers **604** of the other server group **606** (e.g., the server group **606B**). In moving over, the failover manager **616** may redirect or forward communications from the end-consumer devices accessing the application **636** on the initial server group **606** to the second server group **606**. Subsequent to the failover, the end-user consumer devices and the servers **604** of the second server group **606** may exchange communications with each other in accessing the application **636**. In this manner, the failover manager **616** may execute the traffic failover process in accordance with the defined parameters upon invocation from the interaction with the failover execution UI element **634**. The failover manager **616** may automate the various operations involved in failover with minimal or no manual involvement.

[0075] In some embodiments, the failover manager **616** may calculate, determine, or otherwise generate at least one network statistic for each server group **606**. The network statistic may identify or include a measure of performance of the network traffic **638** associated with the application **636** hosted on the servers **604** in the respective server group **606**. The measure of performance may be a single instance or time-series of measurements. The network statistics may include, for example:

latency measuring a delay between the end-user consumer devices with the servers **604** in accessing the application **636**; bandwidth identifying a rate of data exchanged between the end-user consumer devices with the servers **604**; throughput identifying an amount of data successfully communicated between the end-user consumer devices with the servers **604**; jitter corresponding to a variation in latency in the exchanged communications; and an error rate identifying a rate of alterations of the data communicated between the end-user consumer devices with the servers **604** due to network conditions, among others, or any combination thereof. The network statistics may be instrumented by the failover manager **616** (or another computing device). With the generation, the failover manager **616** may store and maintain the network statistics for the server groups **606** on the database **608**. In some embodiments, the servers **604** themselves may generate and store the network statistics as detailed herein on the database **608**. In some embodiments, the failover manager **616** may relay or otherwise provide the network statistics for the server groups **606** to the dashboard handler **612** to present on the user interface **624**.

[0076] In some embodiments, the dashboard handler **612** may provide the user interface **624** to include the network statistics UI element **640**. The network statistics UI element **640** may correspond to a set of graphical user interface (GUI) elements within the user interface **624**, such as a radio button, a check box, a toggle switch, a text box, a window, a prompt, or a container, among others. In some embodiments, the network statistics UI element **640** may be a part of the failover execution UI element **634**. For instance, the GUI elements corresponding to the failover execution UI element **634** may be included in the window including the GUI elements of the network statistics UI element **640**. In some embodiments, the failover execution UI element **634** may be separate from the network statistics UI element **640**. For example, the dashboard handler **312** may present the network statistics UI element **640** in a top portion of the window and the failover execution UI element **634** on a bottom portion of the window.

[0077] The network statistics UI element **640** may include, identify, or otherwise provide information relevant to the traffic failover process, such as the network statistics. The dashboard handler **612** may retrieve, obtain, or otherwise identify the network statistics from the failover manager **616** or another data source (e.g., the database **608**). The network statistics UI element **640** may include or identify the network statistics by application **636**, server **604**, or server group **606**. For example, the network statistics UI element **640** may identify the latency, bandwidth, throughput, the jitter, the error rate, or a combined score for each application **636**, server **604**, or server group **606**. In some embodiments, the dashboard handler **612** may present or provide an indicator for each network statistic. The indication may be, for instance, an enumeration identifier or a color code identifying whether the network statistic for a given application **636** is excellent, good, fair, or poor, among others. By displaying the network statistics in the network statistics UI element **640** in a digestible manner, a user (e.g., the network administrator) can determine whether to invoke the traffic failover process.

[0078] FIG. 7 depicts a screenshot of a network health interface **700** in the dashboard user interface for the system for failover management. The network health interface **700** may identify network traffic statistics categorized by application and server groups. In the depicted example, generally along the top, the network health interface **700** may include a set of application name elements **702A-D** identifying a type of application. For each application, the network health interface **700** may include a set of server group elements **704A-D** and **704'A-D**, each of which may identify network statistics for the given server group and application (e.g., using color coding). Generally along the bottom, the network health interface **700** may include a set of elements **706A-C** identifying network statistics by applications. The network health interface **700** may be used by the network administrator to make decisions regarding whether to invoke network traffic failover process from one server group to another server group for a given application.

[0079] FIGS. 8A and 8B depict screenshots of an execution interface **800** in the dashboard user interface for the system for failover management. Starting with FIG. 8A, in the depicted example,

the execution interface **800** may include an authentication interface **802** to initiate the traffic failover process. The authentication interface **802** may be used to enter a one-time password (OTP) to validate the network administrator prior to invoking the traffic failover process. Moving onto FIG. **8B**, the execution interface **800** may a prompt **810** to enter the application patch identifier to select the application whose network traffic is to be transferred from one server group to another server group. The execution interface **800** may be used by the network administrator to carry out the traffic failover process.

[0080] FIG. **9** depicts a screenshot of a traffic pattern interface **900** in the dashboard user interface for the system for failover management. In the depicted example, the traffic pattern interface **900** may provide additional network statistics in relation to the network traffic for a given application categorized by host server groups (e.g., “GTDC” and “SWDC”) over a given time window. The traffic pattern interface **900** may be used by the network administrator to decide whether to invoke the network traffic failover process from one server group to another server group for a given application.

[0081] FIG. **10** depicts a block diagram of a system **1000** for performance analytics in networked environments. The system **1000** may include at least one automated management service **1002**, one or more servers **1004A-X** (hereinafter generally referred to as servers **1004**) in at least one server group **1006**, and at least one database **1008**, among others, communicatively coupled with one another via at least one network **1010**. The automated management service **1002** may include at least one dashboard handler **1012**, at least one analytics evaluator **1018**, and at least one analytics model **1020**, among others, and may provide at least one user interface **1024**. The database **1008** may store, maintain, or otherwise include historical data **1040**. The user interface **1024** may include at least one analytics retrieval UI element **1030** and a set of analytics results UI elements **1044A-N** (hereinafter generally referred to as analytics results UI elements **1044**), among others. At least one of the servers **1004** in at least one server group **1006** may host resources for at least one application **1036**. The application **1036** may include a set of functions **1038A-N** (hereinafter generally referred to as functions **1038**). Embodiments may comprise additional or alternative components or omit certain components from those of FIG. **10**, and still fall within the scope of this disclosure.

[0082] The dashboard handler **1012** executing on the automated management service **1002** may provide the user interface **1024** including the analytics retrieval UI element **1030**. The analytics retrieval UI element **1030** may be rendered, displayed, or otherwise presented on at least one display of the automated management service **1002** or communicatively coupled with the automated management service **1002**. The analytics retrieval UI element **1030** may correspond to at least one graphical user interface (GUI) element within the user interface **1024**, such as a command button, a slider, a toggle switch, an image, a window, a prompt, or a container, among others, or any combination thereof. The user interface **1024** may be in the manner depicted, for example, in FIGS. **11** and **12**, among others. The analytics retrieval UI element **1030** may provide performance indicators of the functions **1038** supported or provided by the application **1036**. In response to detecting an interaction with the analytics retrieval UI element **1030**, the dashboard handler **1012** may call, invoke, or otherwise execute the analytics evaluator **1018**.

[0083] The analytics evaluator **1018** executing on the automated management service **1002** may calculate, determine, or otherwise generate a set of performance indicators for the corresponding set of functions **1038** of the application **1036**, in response to the interaction with the analytics retrieval UI element **1030**. The generation of the performance indicators may be based on historical data **1040** for the application **1036**. The historical data **1040** may be stored and maintained on a storage (e.g., the database **1008**) using instrumentation of the instance of the application **1036** on one or more of the servers **1004** across server groups **1006**. For each function **1038** of the application **1036**, the historical data **1040** may identify or include, for example: consumption of computing resources (e.g., processor or memory); a number of invocations (or requests); latency

between requests and outputs; down time; success rate in carrying out; and an a number of errors or failures from performing the function **1038**, among others, or any combination thereof. The historical data **1040** may be instrumented or measured on a rolling basis, with overlapping sampling intervals.

[0084] In generating, the analytics evaluator **1018** may calculate, generate, or otherwise determine a performance metric for each function **1038** of the application **1036** based on at least a portion of the historical data **1040**. Each performance metric may identify a respective predicted likelihood that the application **1036** will execute or carry out a corresponding function **1038**. The function **1038** may include one or more defined operations of the application **1036**, such as account information retrieval in an online banking application, a copy and paste operation in a word processor application, account authentication on a video game, or a loading streaming multimedia on a video player, among others. The portion of the historical data **1040** may, for example, correspond to one or more recent time intervals relative to the present. In some embodiments, the analytics evaluator **1018** may determine the performance metric based on a combination of the portion of the historical data **1040**. The combination may be, for example, a summation, a weighted average, or a formula, among others, to generate the performance metric from the historical data **1040**.

[0085] In conjunction, the analytics evaluator **1018** may calculate, generate, or otherwise determine a threshold value with which to compare against the performance metric. The threshold value may be determined based on at least a portion the historical data **1040**. The portion of the historical data **1040** used to determine the threshold value may include more sampling time intervals than the portion of the historical data **1040** used to determine the performance metric. For example, the portion of the historical data **1040** used for the threshold value may correspond to the most recent week, whereas the portion of the historical data **1040** used to determine the performance metric may correspond to the most recent six hours. In some embodiments, the analytics evaluator **1018** may determine the threshold value based on a combination of the portion of the historical data **1040**. The combination may be, for example, a moving average (e.g., weighted or exponential), a weighted sum, or a formula, among others.

[0086] With the determinations, the analytics evaluator **1018** may determine, identify, or otherwise select a performance indicator for each function **1038** of the application **1036**. The performance indicator may be correlated with or identify the predicted likelihood of the application **1036** of successfully performing the function **1038**. The performance indicator may be selected from a positive (or normal) performance indicator corresponding to the performance metric for the function **1038** satisfying (e.g., greater than or equal to) the associated threshold or a negative (or anomalous) performance indicator corresponding to the performance metric for the function **1038** not satisfying (e.g., less than) the associated threshold. To select, the analytics evaluator **1018** may compare the performance metric for the function **1038** with the corresponding threshold. If the performance metric for the function **1038** satisfies the threshold, the analytics evaluator **1018** may select the positive performance indicator. Otherwise, if the performance metric for the function **1038** does not satisfy the threshold, the analytics evaluator **1018** may select the negative performance indicator. Upon the determination, the analytics evaluator **1018** may provide the performance indicators and corresponding performance metrics for the functions **1038** of the application **1036** to the user interface **1024**.

[0087] In some embodiments, the analytics evaluator **1018** may use the analytics model **1020** to determine the set of performance indicators for the set of functions **1038** of the application **1036**. The analytics model **1020** may be, for example, a machine learning (ML) model to process historical data to output performance indicators. The architecture or algorithm used to implement the analytics model **1020** may include, for example, an artificial neural network (ANN), a clustering model (e.g., k nearest neighbors), a regression model (e.g., linear or logistic regression), a random forest, a Bayesian classifier, or a support vector machine (SVM), among others. In

general, the analytics model **1020** may include: a set of inputs corresponding to at least a portion (e.g., the most recent time interval) of the historical data **1040**; and at least one output corresponding to the positive or negative performance indicator; and a set of weights relating the inputs and outputs.

[0088] The analytics model **1020** may be initialized, trained, and established (e.g., by the analytics evaluator **1018** or another computing device) using a training dataset. The training dataset may identify or include the historical data **1040**. The portion of the historical data **1040** used to train the analytics model **1020** may include sampling time intervals prior to the portion of the historical data **1040** to be fed into the analytics model **1020**. For instance, the portion of the historical data **1040** may be from the previous two to five weeks of instrumentation, relative to the most recent sampling. For the training dataset, the input may correspond to the consumption of computing resources; a number of invocations; latency; and down time, among others included in one or more sampling intervals the historical data **1040**. The expected outputs may include the positive or negative performance indicator for the corresponding sampling interval. The weights of the analytics model **1020** may be trained in accordance with supervised learning using the training data. The analytics model **1020** may be continuously trained with updated historical data **1040** from one or more previous intervals of time on a rolling basis.

[0089] With the establishment, the analytics evaluator **1018** may apply the portion of the portion (e.g., the most recent time intervals) of the historical data **1040** for each function **1036** of the application **1038** into the analytics model **1020**. In applying, the analytics evaluator **1018** may feed the historical data **1040** as input into the analytics model **1020**, and may process the historical data **1040** in accordance with the set of weights of the analytics model **1020**. From processing, the analytics evaluator **1018** may produce or generate the performance indicator for the corresponding function **1038**, output from the analytics model **1020**. The analytics evaluator **1018** may repeat the applying of the portions of the historical data **1040** over the set of functions **1038** of the application **1036**. With the determination of the performance indicators, the analytics evaluator **1018** may provide the performance indicators for the functions **1038** of the application **1036** to the user interface **1024**.

[0090] The dashboard handler **1012** may provide the user interface **1024** to include the analytics results UI elements **1042**. The analytics results UI elements **1042** may correspond to a set of graphical user interface (GUI) elements within the user interface **1024**, such as a radio button, a slider, a check box, a toggle switch, a text box, an image, a window, a prompt, or a container, among others. In some embodiments, the analytics results UI elements **1042** may be a part of the analytics retrieval UI element **1030**. For example, the analytics results UI elements **1042** may be included as part of the window of the analytics retrieval UI element **1030**. In some embodiments, the analytics results UI elements **1042** may be separate from the analytics retrieval UI element **1030**. For instance, the analytics retrieval UI element **1030** may reside on a main webpage, and the analytics results UI elements **1042** may be presented on a separate webpage upon interaction with the analytics retrieval UI element **1030**.

[0091] Each analytics result UI element **1042** may correspond to a corresponding performance indicator of the respective function **1038** supported or provided by the application **1036**. The indication may be, for instance, an enumeration identifier or a color code identifying whether the performance indicator for the respective function **1038** is positive (e.g., normal) or negative (e.g., anomalous), among others. For instance, the first analytics result UI element **1042A** may correspond to an account information retrieval feature of an online banking application, and may have a green color to indicate that the feature is operating properly. The second analytics result UI element **1042B** may correspond to a transaction feature of the online banking application, and may have a red color to indicate that the feature is non-operational or otherwise behaving abnormally. The displaying of the performance indicators in the set of analytics results UI elements **1042** in a digestible manner may allow a user (e.g., the network administrator) to diagnose any issues if any

with the application **1036**, the servers **1004** hosting the application **1036**, or the server group **1006**, among others.

[0092] In some embodiments, the dashboard handler **1012** may support or provide a drill-down feature for the performance indicator, upon interaction with at least one of the set of analytics results UI elements **1042**. In response to the interaction with an analytics results UI element **1042**, the dashboard handler **1012** may provide the performance metrics for the performance indication of the corresponding function **1038** of the application **1036**. As discussed above, the performance metric may identify the predicted likelihood of success for the given function **1038**. In some embodiments, in response to the interaction, the dashboard handler **1012** may provide at least a portion of the historical data **1040** for the corresponding function **1038**. The portion of the historical data **1040** may identify of include metrics, such as consumption of computing resources, number of requests, latency, success rate, or number or rate of errors, among others, as discussed above.

[0093] FIG. **11** depicts a screenshot of a performance indicator interface **1100** in the dashboard user interface for the system for performance analytics. The performance indicator interface **1100** may present performance indicators for various functions across multiple applications. In the depicted example, generally in the middle, the performance indicator interface **1100** may a set of lists **1102-C** for each application platform (or operating systems). For each platform, the performance indicator interface **1100** may a set of functions **1110A-N** identifying performance indicators (e.g., using color code indications). The performance indicator interface **1100** may be used by the network administrator to pinpoint certain functions (or transactions) as a cause of issues in a given application.

[0094] FIG. **12** depicts a screenshot of a drill down interface **1200** in the dashboard user interface for the system for performance analytics. The drill down interface **1200** may present additional metrics for a given function of a particular application. In the depicted example, along the top, the drill down interface **1200** may include an indicator element **1205** identifying performance indicators (e.g., using color code) for the identified function over a given time interval. In addition, generally in the middle, the drill down interface **1200** may include a set of graph elements **1210A-D** identifying performance metrics (e.g., failure count trends, slow call count trend, calls per minute, and success count trend) over time. The drill down interface **1200** may be used by the network administrator to view various types of metrics for a particular function in a selected application.

[0095] FIG. **13** depicts a block diagram of a system **1300** for site reliability evaluation in networked environments. The system **1300** may include at least one automated management service **1302**, one or more servers **1304A-X** (hereinafter generally referred to as servers **1304**) in at least one server group **1306**, and at least one database **1308**, among others, communicatively coupled with one another via at least one network **1310**. The automated management service **1302** may include at least one dashboard handler **1312** and at least one reliability evaluator **1320**, among others, and may provide at least one user interface **1322**. The user interface **1324** may include at least one analytics retrieval UI element **1332**, among others. At least one of the servers **1304** in at least one server group **1306** may host resources for at least one application **1336**. Embodiments may comprise additional or alternative components or omit certain components from those of FIG. **13**, and still fall within the scope of this disclosure.

[0096] The dashboard handler **1312** executing on the automated management service **1302** may provide the user interface **1324** including the reliability retrieval UI element **1332**. The reliability retrieval UI element **1332** may be rendered, displayed, or otherwise presented on at least one display of the automated management service **1302** or communicatively coupled with the automated management service **1302**. The reliability retrieval UI element **1332** may correspond to at least one graphical user interface (GUI) element within the user interface **1324**, such as a command button, a slider, a toggle switch, an image, a window, a prompt, or a container, among

others, or any combination thereof. The reliability retrieval UI element **1332** may provide reliability indicators for the application **1036**, the servers **1304**, or the server group **1306**. In response to detecting an interaction with the reliability retrieval UI element **1332**, the dashboard handler **1312** may call, invoke, or otherwise execute the reliability evaluator **1320**.

[0097] The reliability evaluator **1320** executing on the automated management service **1302** may calculate, identify, or determine a set of reliability measures for the application **1036**, the servers **1304**, or the server group **1306**. The reliability evaluator **1320** may retrieve, obtain, or otherwise identify historical data instrumenting the application **1036**, the servers **1304**, or the server group **1306** from the database **1308**. In conjunction, the reliability evaluator **1320** may retrieve, obtain, or otherwise identify one or more service level objectives (SLOs) for the application **1036**, the servers **1304**, or the server group **1306**. With the identification, the reliability evaluator **1320** may compare the historical data with the SLOs. Based on the determination, the reliability evaluator **1320** may determine the reliability measures. With the determination, the reliability evaluator **1320** may provide the reliability indicators for presentation on the user interface **1324**, in the manner depicted in FIGS. **14-17**.

[0098] FIG. **14** depicts a screenshot of an on-boarding interface **1400** in the dashboard user interface for the system for site reliability evaluation. The on-boarding interface **1400** may be a graphical user interface used to enter various parameters for SLOs. FIG. **15** depicts a screenshot of a reliability indication interface **1500** in the dashboard user interface for the system for site reliability evaluation. The reliability indication interface **1500** may present various statistics relevant to whether the SLOs are being met. FIG. **16** depicts a screenshot of a drill down interface **1600** in the dashboard user interface for the system for site reliability evaluation. The drill down interface **1600** may provide additional relevant statistics for a particular application, servers, or server groups. FIG. **17** depicts a screenshot of a heat map interface **1700** in the dashboard user interface for the system for site reliability evaluation. The heat map interface **1700** may present a set of reliability indicators for particular applications or platforms over multiple time intervals.

[0099] FIG. **18** depicts a block diagram of a system **1800** for services assistance in a system for managing networked environments. The system **1800** may include at least one automated management service **1802**, one or more servers **1804A-X** (hereinafter generally referred to as servers **1804**) in at least one server group **1806**, and at least one database **1808**, among others, communicatively coupled with one another via at least one network **1810**. The automated management service **1802** may include at least one dashboard handler **1812** and at least one assistant handler **1822**, among others, and may provide at least one user interface **1824**. The user interface **1824** may include at least one assistant invocation UI element **1834**, among others. At least one of the servers **1804** in at least one server group **1806** may host resources for at least one application **1836**. Embodiments may comprise additional or alternative components or omit certain components from those of FIG. **18**, and still fall within the scope of this disclosure.

[0100] The dashboard handler **1812** executing on the automated management service **1802** may provide the user interface **1824** including the assistant invocation UI element **1834**. The assistant invocation UI element **1834** may be rendered, displayed, or otherwise presented on at least one display of the automated management service **1002** or communicatively coupled with the automated management service **1802**. The assistant invocation UI element **1834** may correspond to at least one graphical user interface (GUI) element within the user interface **1824**, such as a command button, a slider, a toggle switch, an image, a window, a prompt, or a container, among others, or any combination thereof. The assistant invocation UI element **1834** may provide an interface to a customer agent service to obtain health checks or enter incidents for the application **1836**, the servers **1804**, or the server group **1806**. In response to detecting an interaction with the assistant invocation UI element **1834**, the dashboard handler **1812** may call, invoke, or otherwise execute the assistant handler **1822**.

[0101] The assistant handler **1822** executing on the automated management service **1802** may

retrieve, obtain, or identify a health status of the application **1836**, the servers **1804**, or the server group **1806** selected via the user interface **1824**. The assistant handler **1822** may invoke a customer services agent (e.g., a digital assistant application) to retrieve the health status of the indicated the application **1836**, the servers **1804**, or the server group **1806**. The assistant handler **1822** may also provide invoke the customer services agent to enter details regarding an incident (e.g., an outage or an interruption) of the application **1836**, the servers **1804**, or the server group **1806**. The customer services agent may handle incident ticketing to prevent duplicates or aggregate similar incidents. The input and output interfaces for the customer service agent may be presented in the interface **1824** in the manner depicted in FIGS. **19** and **20**.

[0102] FIG. **19** depicts a screenshot of a health check interface **1900** in the dashboard user interface for the system for services assistance. In the depicted example, the health check interface **1900** may present a health status of various functions of an application. The health check interface **1900** may be also used to enter an impacted application to report incidents. FIG. **20** depicts a screenshot of a query interface **2000** in the dashboard user interface for the system for services assistance. In the depicted example, the query interface **2000** may enter additional information for the incident. The query interface **2000** may be used to submit incident reports to the customer service agent.

[0103] FIG. **21** depicts a flow diagram of a method of managing networked environments. Embodiments may include additional, fewer, or different operations from those described in the method **2100**. The method **2100** may be performed by a service (e.g., an automated management service) executing machine-readable software code, though it should be appreciated that the various operations may be performed by one or more computing devices and/or processors. At step **2105**, the service may provide a dashboard user interface. The dashboard user interface may include a first element to invoke patch management, a second element to execute traffic failover, a third element to retrieve predictive analytics, a fourth element to provide reliability indications, and a fifth element to invoke a services assistant.

[0104] At step **2110**, the service may monitor for an interaction with one of the elements of the dashboard user interface. At step **2115**, if the interaction is detected on the dashboard user interface the service may determine which process to invoke. At step **2120**, if the interaction is with the first element, the service may execute the patch process to shut down a server group, install the patch, and perform validation. At step **2125**, if the interaction is with the second element, the service may perform a failover process to transfer network traffic for an identifier application from one group of servers to another group of services.

[0105] In addition, at step **2130**, if the interaction is with the third element, the service may provide analytics by generating performance indicators for functions of the application. At step **2135**, if the interaction is with the fourth element, the service may provide reliability indicators in accordance with service level objections (SLOs). At step **2140**, if the interaction is with the fifth element, the service may invoke the customer services agent to check health statuses of applications and managing ticketing of incidents. At step **2145**, the service may provide the output from the performed process on the dashboard user interface. Subsequently, the service may repeat the method **2100** from the step **2110**.

[0106] The foregoing method descriptions and the process flow diagrams are provided merely as illustrative examples and are not intended to require or imply that the steps of the various embodiments must be performed in the order presented. The steps in the foregoing embodiments may be performed in any order. Words such as “then,” “next,” etc. are not intended to limit the order of the steps; these words are simply used to guide the reader through the description of the methods. Although process flow diagrams may describe the operations as a sequential process, many of the operations can be performed in parallel or concurrently. In addition, the order of the operations may be re-arranged. A process may correspond to a method, a function, a procedure, a subroutine, a subprogram, and the like. When a process corresponds to a function, the process termination may correspond to a return of the function to a calling function or a main function.

[0107] The various illustrative logical blocks, modules, circuits, and algorithm steps described in connection with the embodiments disclosed herein may be implemented as electronic hardware, computer software, or combinations of both. To clearly illustrate this interchangeability of hardware and software, various illustrative components, blocks, modules, circuits, and steps have been described above generally in terms of their functionality. Whether such functionality is implemented as hardware or software depends upon the particular application and design constraints imposed on the overall system. Skilled artisans may implement the described functionality in varying ways for each particular application, but such implementation decisions should not be interpreted as causing a departure from the scope of the present disclosure.

[0108] Embodiments implemented in computer software may be implemented in software, firmware, middleware, microcode, hardware description languages, or any combination thereof. A code segment or machine-executable instructions may represent a procedure, a function, a subprogram, a program, a routine, a subroutine, a module, a software package, a class, or any combination of instructions, data structures, or program statements. A code segment may be coupled to another code segment or a hardware circuit by passing and/or receiving information, data, arguments, parameters, or memory contents. Information, arguments, parameters, data, etc. may be passed, forwarded, or transmitted via any suitable means including memory sharing, message passing, token passing, network transmission, etc.

[0109] The actual software code or specialized control hardware used to implement these systems and methods is not limiting. Thus, the operation and behavior of the systems and methods were described without reference to the specific software code being understood that software and control hardware can be designed to implement the systems and methods based on the description herein.

[0110] When implemented in software, the functions may be stored as one or more instructions or code on a non-transitory computer-readable or processor-readable storage medium. The steps of a method or algorithm disclosed herein may be embodied in a processor-executable software module, which may reside on a computer-readable or processor-readable storage medium. A non-transitory computer-readable or processor-readable media includes both computer storage media and tangible storage media that facilitate transfer of a computer program from one place to another. A non-transitory processor-readable storage media may be any available media that may be accessed by a computer. By way of example, and not limitation, such non-transitory processor-readable media may comprise RAM, ROM, EEPROM, CD-ROM or other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other tangible storage medium that may be used to store desired program code in the form of instructions or data structures and that may be accessed by a computer or processor. Disk and disc, as used herein, include compact disc (CD), laser disc, optical disc, digital versatile disc (DVD), floppy disk, and Blu-ray disc where disks usually reproduce data magnetically, while discs reproduce data optically with lasers.

Combinations of the above should also be included within the scope of computer-readable media. Additionally, the operations of a method or algorithm may reside as one or any combination or set of codes and/or instructions on a non-transitory processor-readable medium and/or computer-readable medium, which may be incorporated into a computer program product.

[0111] The preceding description of the disclosed embodiments is provided to enable any person skilled in the art to make or use the present disclosure. Various modifications to these embodiments will be readily apparent to those skilled in the art, and the generic principles defined herein may be applied to other embodiments without departing from the spirit or scope of the disclosure. Thus, the present disclosure is not intended to be limited to the embodiments shown herein but is to be accorded the widest scope consistent with the following claims and the principles and novel features disclosed herein.

[0112] While various aspects and embodiments have been disclosed, other aspects and embodiments are contemplated. The various aspects and embodiments disclosed are for purposes

of illustration and are not intended to be limiting, with the true scope and spirit being indicated by the following claims.

Claims

1. A method of configuring server groups, comprising: providing, by one or more processors, for display a user interface for configuring a plurality of groups of servers, the user interface comprising a plurality of elements to execute a corresponding plurality of processes, the plurality of elements including: a first element configured to invoke, responsive to a first interaction, a first process to provide predictive analytics by providing a performance indicator with a function based on first historical data for the function; and a second element configured to invoke, responsive to a second interaction, a second process to provide a reliability metric for an application on at least one of the plurality of groups of servers determined based on second historical data of the application in accordance with a service level objective (SLO) for the application; identifying, by the one or more processors, from the plurality of elements, an element upon which an interaction is detected; and executing, by the one or more processors, a process of the plurality of processes based on the identified element.
2. The method of claim 1, wherein the plurality of elements of the user interface further comprises a third element configured to invoke, responsive to a third interaction, a third process to transfer network traffic from a first group of servers to a second group of servers of the plurality of groups of servers, and wherein the third process further comprises identifying the second group of servers as available to host resources for an application provided by the first group of servers and transferring the network traffic for the application from the first group servers to the second group of servers.
3. The method of claim 1, wherein the plurality of elements of the user interface further comprises a third element configured to invoke, responsive to a third interaction, a third process to provide a digital assistant application to identify a status of at least one group of the plurality of groups of servers, and wherein the third process further comprises providing, for display via the user interface, a user interface element of the digital assistant application to accept input regarding the status of the at least one group.
4. The method of claim 1, wherein the plurality of elements of the user interface further comprises a third element configured to invoke, responsive to a third interaction, a third process to install a patch on at least one group of servers of the plurality of groups of servers, and wherein the third process further comprises shutting down the at least one group of servers, installing the patch on the at least one group of servers, validating the installation of the patch, and providing an identification of a status of the installation as output via the user interface.
5. The method of claim 1, wherein the first process further comprises: generating a performance metric identifying a predicted likelihood of the application successfully carrying out the function based on the first historical data of the function on at least one of the plurality of groups of servers; and selecting from a plurality of performance indicators, the performance indicator based on the performance metric.
6. The method of claim 5, wherein generating the performance metric further comprises applying a machine learning (ML) model on the first historical data to generate the performance metric.
7. The method of claim 1, wherein the second process further comprises: comparing the second historical data of the application on at least one group of servers of the plurality of groups of servers, with the SLO for the application; and generating the reliability metric for the application in accordance with comparison of the second historical data with the SLO.
8. The method of claim 7, wherein generating the reliability metric further comprises generating the reliability metric indicating a number of SLO breaches for the application on at least one group of servers.

9. The method of claim 1, further comprising generating, by the one or more processors, for display via the user interface, an output in accordance with execution of the process.

10. The method of claim 1, wherein the plurality of groups of servers comprises at least one of a first group of servers corresponding to a data center, a second group of servers corresponding to a branch office, or a third group of servers corresponding to a site.

11. A system for configuring server groups, comprising: one or more processors coupled with memory, configured to: provide, for display a user interface for configuring a plurality of groups of servers, the user interface comprising a plurality of elements to execute a corresponding plurality of processes, the plurality of elements including: a first element configured to invoke, responsive to a first interaction, a first process to provide predictive analytics by providing a performance indicator with a function based on first historical data for the function; and a second element configured to invoke, responsive to a second interaction, a second process to provide a reliability metric for an application on at least one of the plurality of groups of servers determined based on second historical data of the application in accordance with a service level objective (SLO) for the application; identify, from the plurality of elements, an element upon which an interaction is detected; and execute a process of the plurality of processes based on the identified element.

12. The system of claim 11, wherein the plurality of elements of the user interface further comprises a third element configured to invoke, responsive to a third interaction, a third process to transfer network traffic from a first group of servers to a second group of servers of the plurality of groups of servers, and wherein the third process further comprises identifying the second group of servers as available to host resources for an application provided by the first group of servers and transferring the network traffic for the application from the first group servers to the second group of servers.

13. The system of claim 11, wherein the plurality of elements of the user interface further comprises a third element configured to invoke, responsive to a third interaction, a third process to provide a digital assistant application to identify a status of at least one group of the plurality of groups of servers, and wherein the third process further comprises providing, for display via the user interface, a user interface element of the digital assistant application to accept input regarding the status of the at least one group.

14. The system of claim 11, wherein the plurality of elements of the user interface further comprises a third element configured to invoke, responsive to a third interaction, a third process to install a patch on at least one group of servers of the plurality of groups of servers, and wherein the third process further comprises shutting down the at least one group of servers, installing the patch on the at least one group of servers, validating the installation of the patch, and providing an identification of a status of the installation as output via the user interface.

15. The system of claim 11, wherein the first process further comprises: generating a performance metric identifying a predicted likelihood of the application successfully carrying out the function based on the first historical data of the function on at least one of the plurality of groups of servers; and selecting from a plurality of performance indicators, the performance indicator based on the performance metric.

16. The system of claim 15, wherein the one or more processors are further configured to apply a machine learning (ML) model on the first historical data to generate the performance metric.

17. The system of claim 11, wherein the second process further comprises: comparing the second historical data of the application on at least one group of servers of the plurality of groups of servers, with the SLO for the application; and generating the reliability metric for the application in accordance with comparison of the second historical data with the SLO.

18. The system of claim 17, wherein the one or more processors are further configured to generate the reliability metric indicating a number of SLO breaches for the application on at least one group of servers.

19. The system of claim 11, wherein the one or more processors are further configured to generate,

for display via the user interface, an output in accordance with execution of the process.

20. The system of claim 11, wherein the plurality of groups of servers comprises at least one of a first group of servers corresponding to a data center, a second group of servers corresponding to a branch office, or a third group of servers corresponding to a site.
