



US 20250260666A1

(19) **United States**

(12) **Patent Application Publication**  
**Rosomakho**

(10) **Pub. No.: US 2025/0260666 A1**

(43) **Pub. Date: Aug. 14, 2025**

(54) **SYSTEMS AND METHODS FOR ENDPOINT  
PROCESS METADATA BASED POLICY  
ENFORCEMENT**

(52) **U.S. CL.**  
CPC ..... **H04L 63/0227** (2013.01); **H04L 63/1425**  
(2013.01); **H04L 63/166** (2013.01)

(71) Applicant: **Zscaler, Inc.**, San Jose, CA (US)

(57) **ABSTRACT**

(72) Inventor: **Yaroslav Rosomakho**, Thatcham (GB)

(73) Assignee: **Zscaler, Inc.**, San Jose, CA (US)

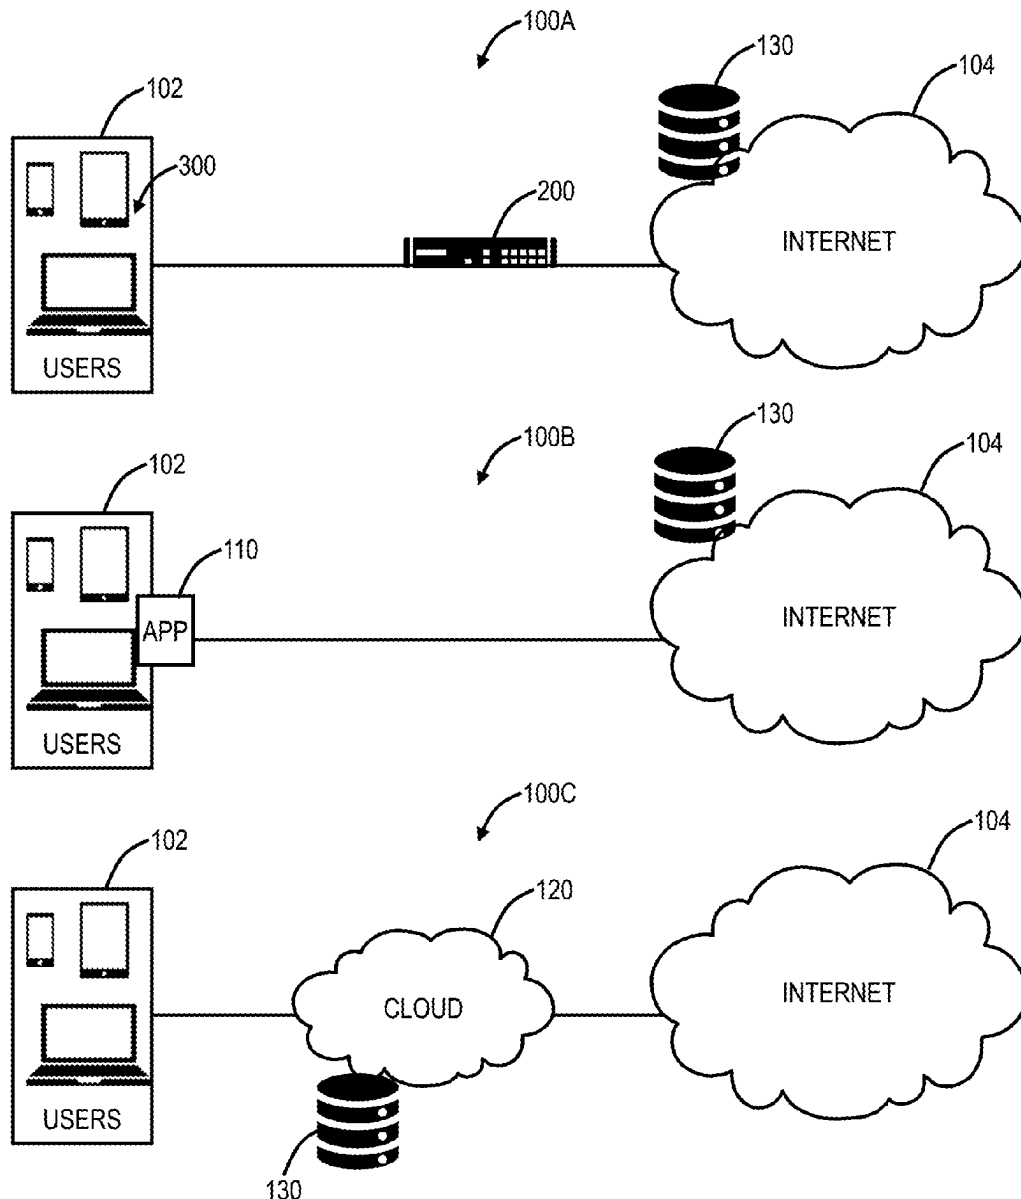
(21) Appl. No.: **18/440,007**

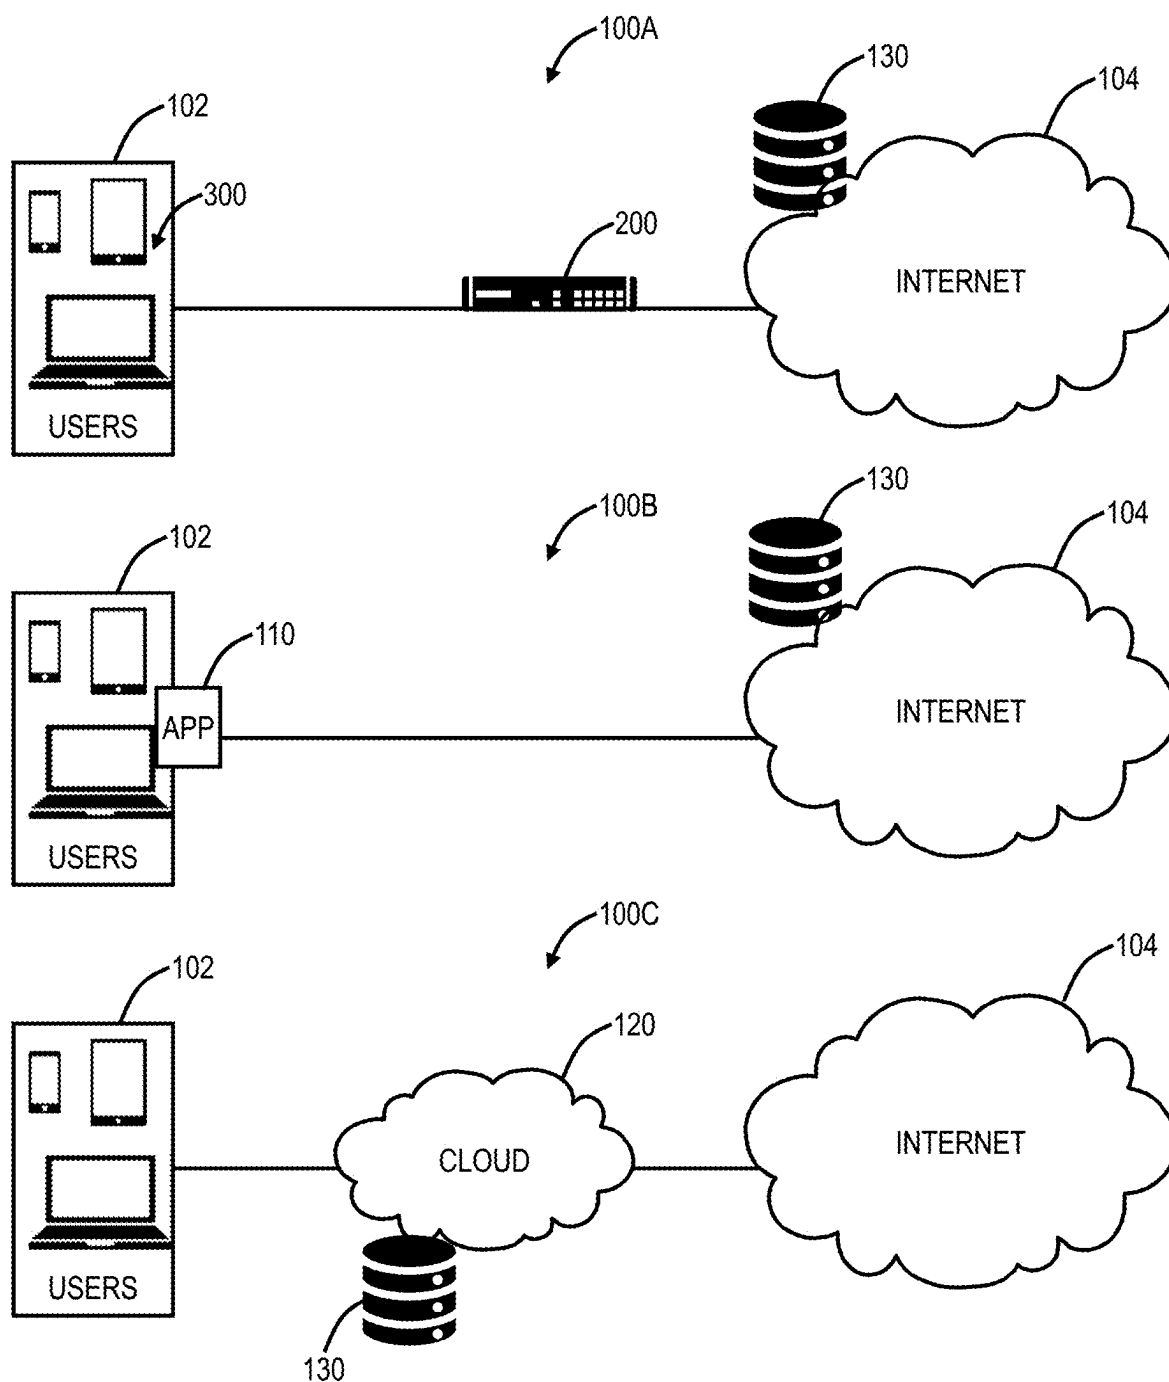
(22) Filed: **Feb. 13, 2024**

**Publication Classification**

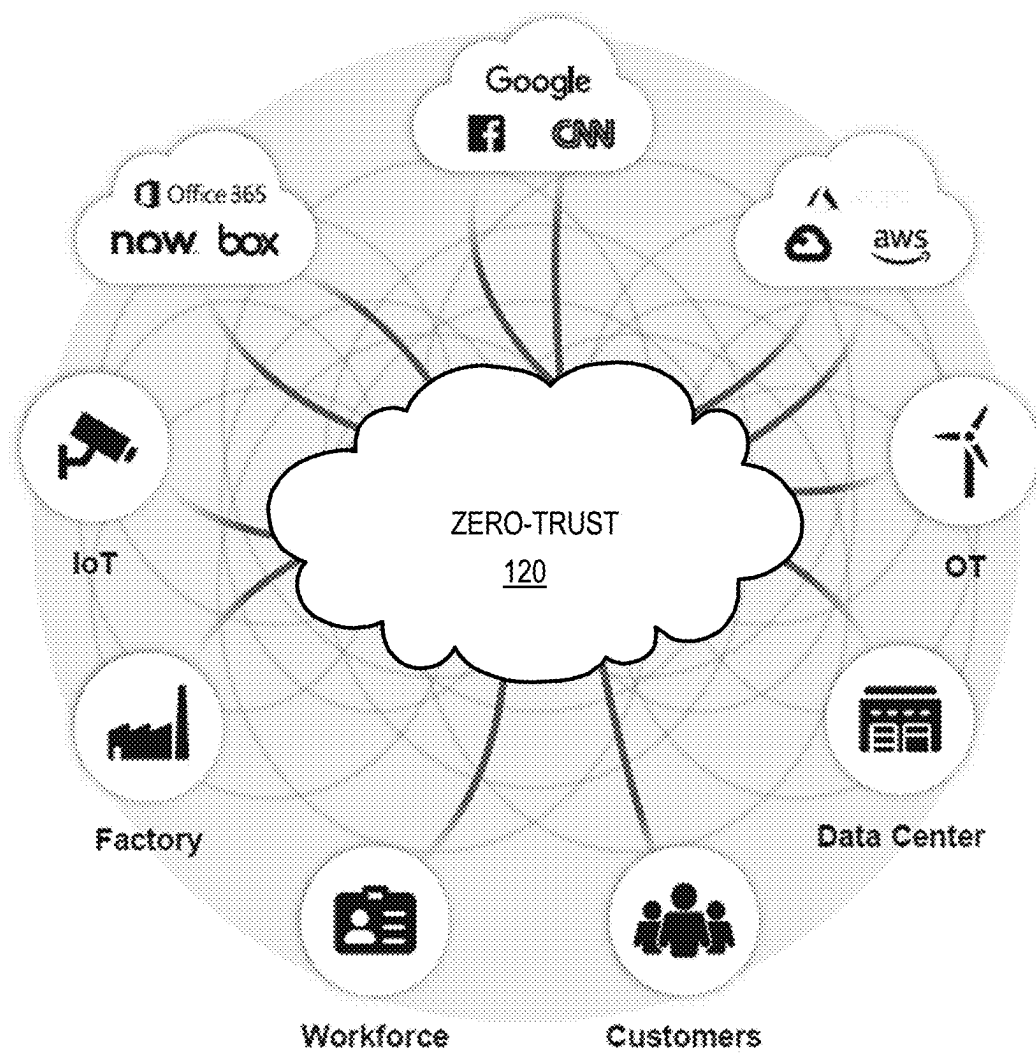
(51) **Int. Cl.**  
**H04L 9/40** (2022.01)

Systems and methods for endpoint application metadata based policy enforcement include monitoring traffic via a cloud, the traffic being monitored inline between one or more endpoints and one or more destinations; identifying, within a request from an endpoint, endpoint process metadata associated with an endpoint process used to make the request; processing the endpoint process metadata; and performing one or more actions on the request based on the processing. The endpoint process metadata can be collected by a connector application executing on the one or more endpoints, and forwarded to the cloud in-band therefrom.



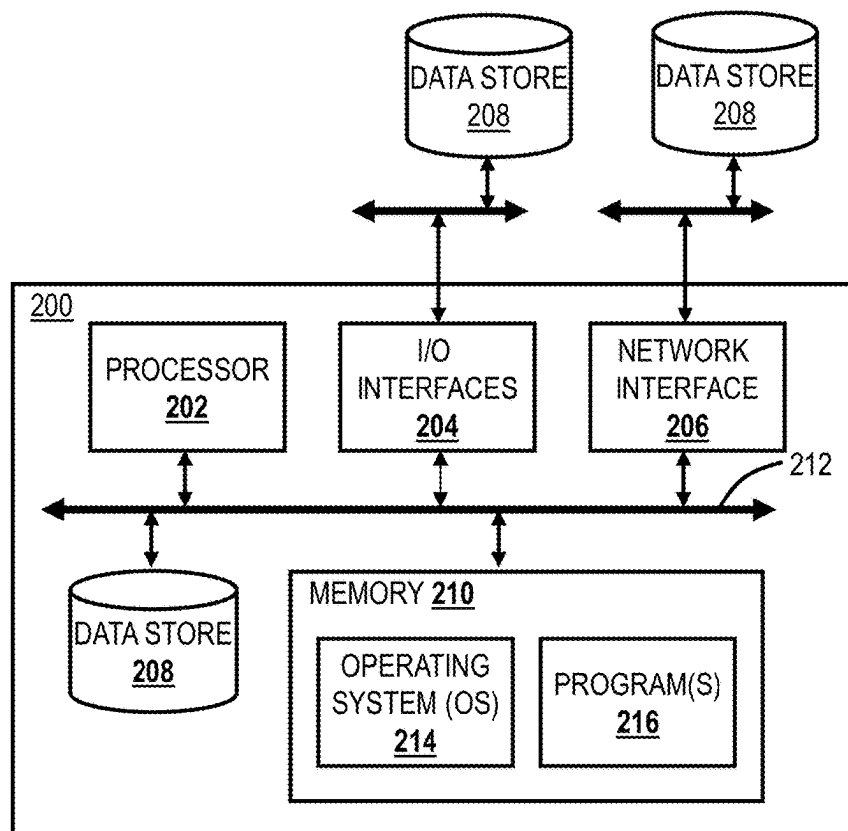


**FIG. 1A**

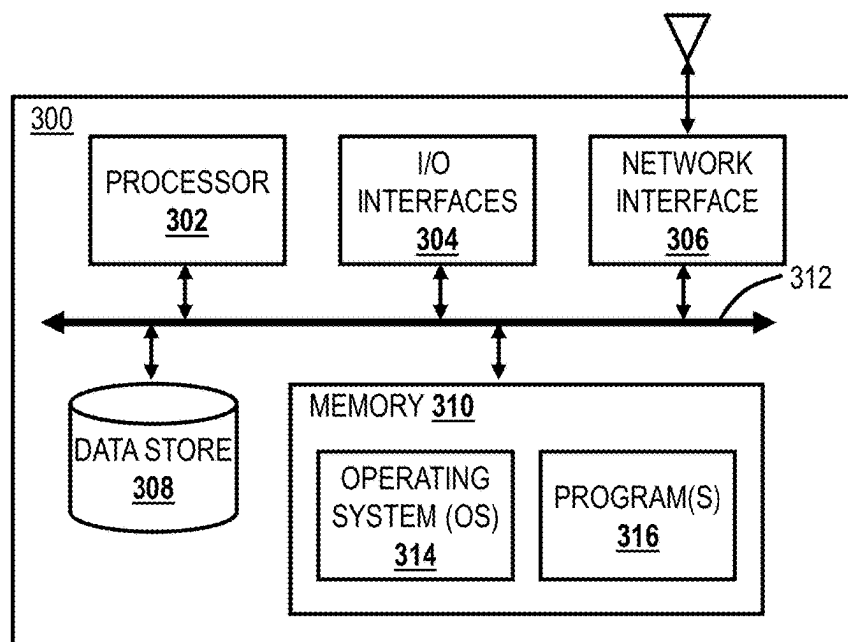


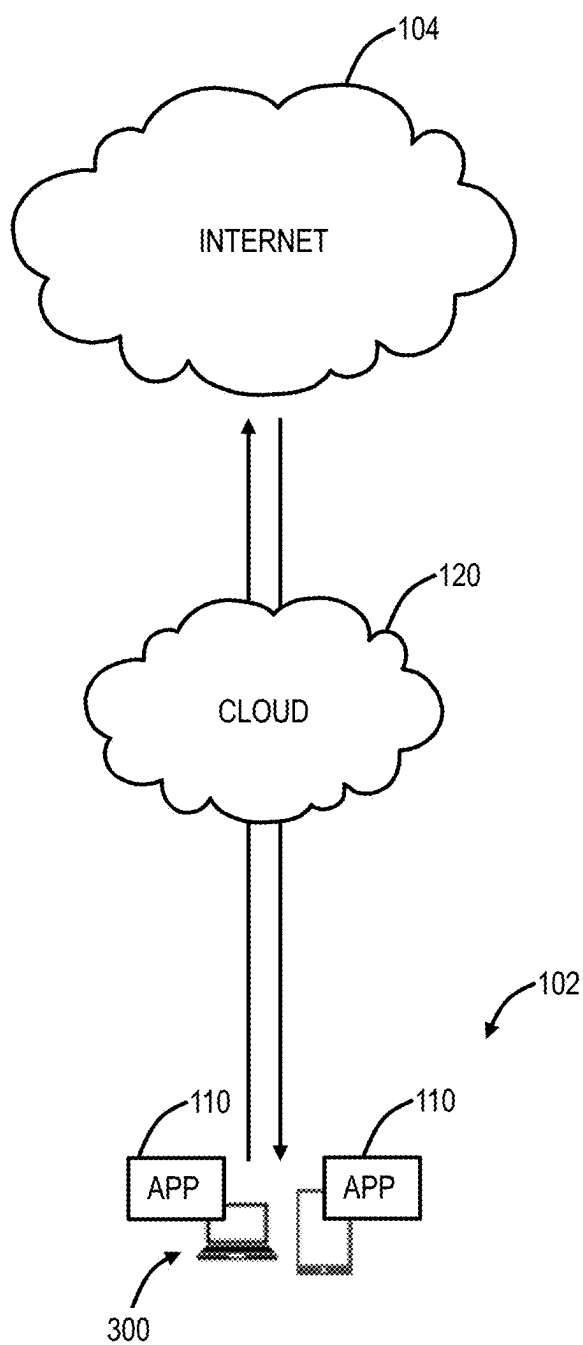
**FIG. 1B**

**FIG. 2**

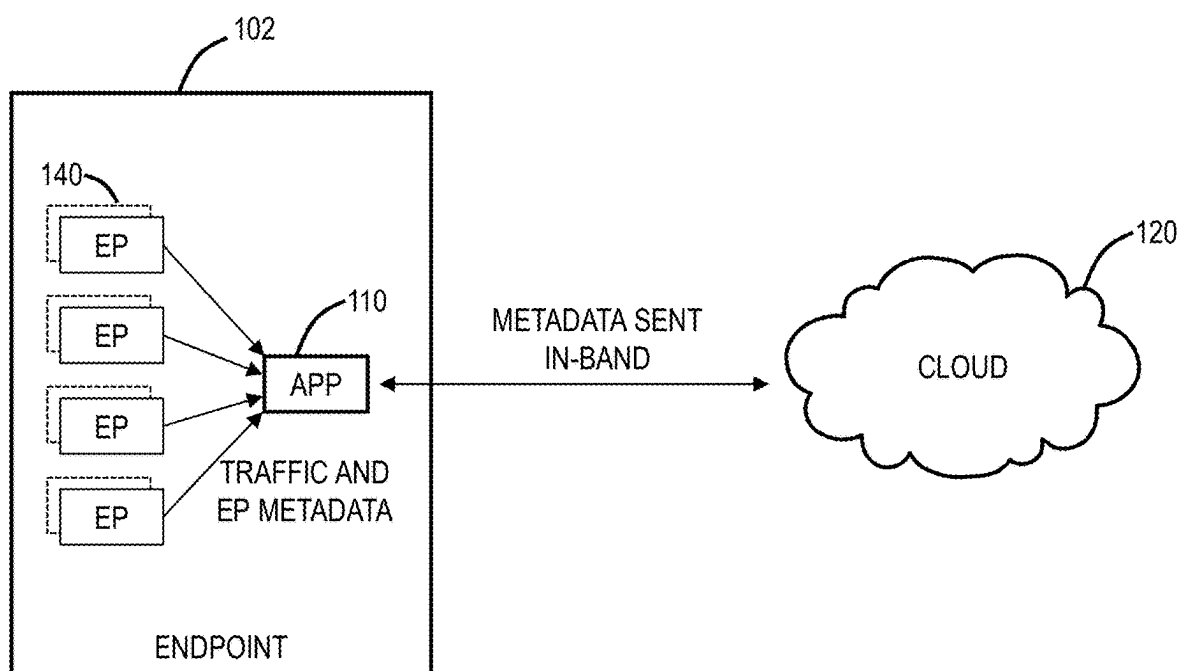


**FIG. 3**

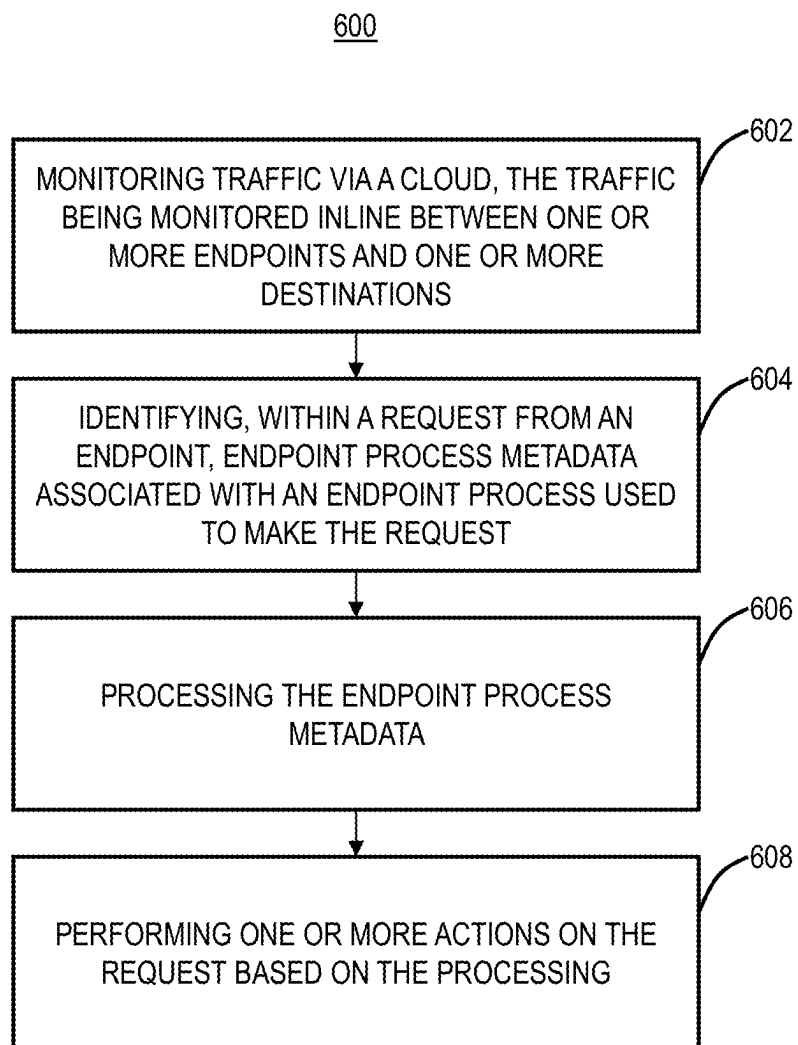




**FIG. 4**



**FIG. 5**

**FIG. 6**

## SYSTEMS AND METHODS FOR ENDPOINT PROCESS METADATA BASED POLICY ENFORCEMENT

### FIELD OF THE DISCLOSURE

[0001] The present disclosure generally relates to network and cloud security. More particularly, the present disclosure relates to systems and methods for endpoint process metadata based policy enforcement.

### BACKGROUND OF THE DISCLOSURE

[0002] The number of endpoints in customer cloud environments is rising. This is due to enterprise networks being cloud-based, having applications, resources, and other destinations moved to the cloud. Managing traffic to and from endpoints and destinations is becoming increasingly comprehensive as enterprises do not have the ability to directly manage every endpoint used to access their data. By not having complete control of software running on the many endpoint devices utilized to access enterprise data and resources, it is difficult to accurately enforce cybersecurity policy. This results in an increased risk for the enterprise data residing on unsecured and unmanaged devices as well as the security risks in access to the Internet. Having visibility into the software responsible for originating traffic flows from endpoint devices enables cybersecurity services to make better decisions at an endpoint process level, which allows for early anomaly detection and new threat protection capabilities. The present disclosure provides systems and methods for collecting and utilizing endpoint process metadata for policy enforcement, and model training purposes.

### BRIEF SUMMARY OF THE DISCLOSURE

[0003] The present disclosure relates to systems and methods for endpoint process metadata based policy enforcement. In various embodiments, the present disclosure includes a method having steps, a processing device configured to implement the steps, a cloud-based system configured to implement the steps, and as a non-transitory computer-readable medium storing instructions for programming one or more processors to execute the steps. The steps include monitoring traffic via a cloud, the traffic being monitored inline between one or more endpoints and one or more destinations; identifying, within a request from an endpoint, endpoint process metadata associated with an endpoint process used to make the request; processing the endpoint process metadata; and performing one or more actions on the request based on the processing.

[0004] The steps can further include wherein the one or more actions include blocking the request, allowing the request, and alerting a user of the endpoint. The request can be forwarded to the cloud via a connector application executing on the endpoint. The connector application can be adapted to collect endpoint process metadata from an endpoint process used to make the request. The connector application can be adapted to forward the endpoint process metadata to the cloud in-band with the request. The endpoint process metadata can include one or more of a signor name, signature timestamp, product name, product version, original file name, executable hash, Process Identity (PID), process name, process owner, process running duration of the endpoint process, and other properties provided by the endpoint operating system. When the endpoint process

metadata includes the product version of the endpoint process, the steps can further include processing the endpoint process metadata to determine if the product version of the endpoint process is the latest version available or an approved version; and performing one or more actions on the request based thereon. Responsive to determining the product version of the endpoint process is not the latest version available or not an approved version, the steps can include blocking the request and alerting a user of the endpoint. The steps can further include logging the endpoint process metadata. The steps can further include providing a customer of the cloud an inventory of all endpoint processes being utilized on their endpoints based on the logging.

### BRIEF DESCRIPTION OF THE DRAWINGS

[0005] The present disclosure is illustrated and described herein with reference to the various drawings, in which like reference numbers are used to denote like system components/method steps, as appropriate, and in which:

[0006] FIG. 1A is a network diagram of three example network configurations of cybersecurity monitoring and protection of a user.

[0007] FIG. 1B is a logical diagram of the cloud operating as a zero-trust platform.

[0008] FIG. 2 is a block diagram of a server.

[0009] FIG. 3 is a block diagram of a computing device.

[0010] FIG. 4 is a diagram of an exemplary network configuration illustrating an application on computing devices configured to operate through the cloud.

[0011] FIG. 5 is a flow diagram of a network configuration including an endpoint, a connector application, and the cloud.

[0012] FIG. 6 is a flow chart of a process for endpoint application metadata based policy enforcement.

### DETAILED DESCRIPTION OF THE DISCLOSURE

[0013] Again, the present disclosure relates to systems and methods for endpoint process metadata based policy enforcement. In various embodiments, the present systems include collecting metadata from one or more endpoint processes such as applications, web browsers, and the like. This metadata is then sent in-band with a request to a cloud for processing. The cloud is adapted to enforce various cybersecurity services by monitoring network traffic inline. Based on processing the endpoint process metadata sent in-band with a request from an endpoint, the cloud can determine whether to allow the request or perform one or more other actions. By employing the systems described herein, the cloud can extend zero trust security down to the endpoint process level. Additionally, the by logging such endpoint process metadata, customers can be provided with accurate inventories of endpoint processes which are used by their endpoints.

#### § 1.0 Cybersecurity Monitoring and Protection Examples

[0014] FIG. 1A is a network diagram of three example network configurations 100A, 100B, 100C of cybersecurity monitoring and protection of an endpoint 102. Those skilled in the art will recognize these are some examples for illustration purposes, there may be other approaches to cybersecurity monitoring (as well as providing generalized



services), and these various approaches can be used in combination with one another as well as individually. Also, while shown for a single endpoint **102**, practical embodiments will handle a large volume of endpoints **102**, including multi-tenancy. In this example, the endpoint **102** communicates on the Internet **104**, including accessing cloud services, Software-as-a-Service, etc. (each may be offered via computing resources, such as, e.g., using one or more servers **200** as illustrated in FIG. 2).

**[0015]** Note, the term endpoint **102** is used herein to refer to any computing device (see FIG. 3 for an example computing device **300**) which can communicate on a network. The endpoint **102** can be associated with a user and include laptops, tablets, mobile phones, desktops, etc. Further, the endpoint can also mean machines, workloads, IoT devices, or simply anything associated with the company that connects to the Internet, a Local Area Network (LAN), etc.

**[0016]** As part of offering cybersecurity through these example network configurations **100A**, **100B**, **100C**, there is a large amount of cybersecurity data obtained. Various embodiments of the present disclosure focus on using this cybersecurity data along with a customer's data to perform various security tasks including developing customer machine learning models and other security platforms of the like.

**[0017]** The network configuration **100A** includes a server **200** located between the endpoint **102** and the Internet **104**. For example, the server **200** can be a proxy, a gateway, a Secure Web Gateway (SWG), Secure Internet and Web Gateway, Secure Access Service Edge (SASE), Secure Service Edge (SSE), Cloud Application Security Broker (CASB), etc. The server **200** is illustrated located inline with the endpoint **102** and configured to monitor the endpoint **102**. In other embodiments, the server **200** does not have to be inline. For example, the server **200** can monitor requests from the endpoint **102** and responses to the endpoint **102** for one or more security purposes, as well as allow, block, warn, and log such requests and responses. The server **200** can be on a local network associated with the endpoint **102** as well as external, such as on the Internet **104**. Also, while described as a server **200**, this can also be a router, switch, appliance, virtual machine, etc. The network configuration **100B** includes an application **110** (connector application) that is executed on the computing device **300**. The connector application **110** can perform similar functionality as the server **200**, as well as coordinated functionality with the server **200** (a combination of the network configurations **100A**, **100B**). Finally, the network configuration **100C** includes a cloud service **120** configured to monitor the endpoint **102** and perform security-as-a-service. Of course, various embodiments are contemplated herein, including combinations of the network configurations **100A**, **100B**, **100C** together.

**[0018]** The cybersecurity monitoring and protection can include firewall, intrusion detection and prevention, Uniform Resource Locator (URL) filtering, content filtering, bandwidth control, Domain Name System (DNS) filtering, protection against advanced threat (malware, spam, Cross-Site Scripting (XSS), phishing, etc.), data protection, sandboxing, antivirus, and any other security technique. Any of these functionalities can be implemented through any of the network configurations **100A**, **100B**, **100C**. A firewall can provide Deep Packet Inspection (DPI) and access controls across various ports and protocols as well as being applica-

tion and user aware. The URL filtering can block, allow, or limit website access based on policy for a user, group of users, or entire organization, including specific destinations or categories of URLs (e.g., gambling, social media, etc.). The bandwidth control can enforce bandwidth policies and prioritize critical applications such as relative to recreational traffic. DNS filtering can control and block DNS requests against known and malicious destinations.

**[0019]** The intrusion prevention and advanced threat protection can deliver full threat protection against malicious content such as browser exploits, scripts, identified botnets and malware callbacks, etc. The sandbox can block zero-day exploits (just identified) by analyzing unknown files for malicious behavior. The antivirus protection can include antivirus, antispyware, antimalware, etc. protection for the endpoints **102**, using signatures sourced and constantly updated. The DNS security can identify and route command-and-control connections to threat detection engines for full content inspection. The DLP can use standard and/or custom dictionaries to continuously monitor the endpoints **102**, including compressed and/or Transport Layer Security (TLS) or Secure Sockets Layer (SSL)-encrypted traffic.

**[0020]** In typical embodiments, the network configurations **100A**, **100B**, **100C** can be multi-tenant and can service a large volume of the endpoints **102**. Newly discovered threats can be promulgated for all tenants practically instantaneously. The endpoints **102** can be associated with a tenant, which may include an enterprise, a corporation, an organization, etc. That is, a tenant is a group of users who share a common grouping with specific privileges, i.e., a unified group under some IT management. The present disclosure can use the terms tenant, enterprise, organization, enterprise, corporation, company, etc. interchangeably and refer to some group of endpoints **102** under management by an IT group, department, administrator, etc., i.e., some group of endpoints **102** that are managed together. One advantage of multi-tenancy is the visibility of cybersecurity threats across a large number of endpoints **102**, across many different organizations, across the globe, etc. This provides a large volume of data to analyze, use machine learning techniques on, develop comparisons, etc. The present disclosure can use the term "service provider" to denote an entity providing the cybersecurity monitoring and a "customer" as a company (or any other grouping of endpoints **102**).

**[0021]** Of course, the cybersecurity techniques above are presented as examples. Those skilled in the art will recognize other techniques are also contemplated herewith. That is, any approach to cybersecurity that can be implemented via any of the network configurations **100A**, **100B**, **100C**. Also, any of the network configurations **100A**, **100B**, **100C** can be multi-tenant with each tenant having its own endpoints **102** and configuration, policy, rules, etc.

### § 1.1 Cloud Monitoring

**[0022]** The cloud **120** can scale cybersecurity monitoring and protection with near-zero latency on the endpoints **102**. Also, the cloud **120** in the network configuration **100C** can be used with or without the connector application **110** in the network configuration **100B** and the server **200** in the network configuration **100A**. Logically, the cloud **120** can be viewed as an overlay network between endpoints **102** and the Internet **104** (and cloud services, SaaS, other destinations, etc.). Previously, the IT deployment model included

enterprise resources and applications stored within a data center (i.e., physical devices) behind a firewall (perimeter), accessible by employees, partners, contractors, etc. on-site or remote via Virtual Private Networks (VPNs), etc. The cloud **120** replaces the conventional deployment model. The cloud **120** can be used to implement these services in the cloud without requiring the physical appliances and management thereof by enterprise IT administrators. As an ever-present overlay network, the cloud **120** can provide the same functions as the physical devices and/or appliances regardless of geography or location of the endpoints **102**, as well as independent of platform, operating system, network access technique, network access provider, etc.

**[0023]** There are various techniques to forward traffic between the endpoints **102** and the cloud **120**. A key aspect of the cloud **120** (as well as the other network configurations **100A**, **100B**) is that all traffic between the endpoints **102** and the Internet **104** is monitored. All of the various monitoring approaches can include log data **130** accessible by a management system, management service, analytics platform, and the like. For illustration purposes, the log data **130** is shown as a data storage element and those skilled in the art will recognize the various compute platforms described herein can have access to the log data **130** for implementing any of the techniques described herein for risk quantification. In an embodiment, the cloud **120** can be used with the log data **130** from any of the network configurations **100A**, **100B**, **100C**, as well as other data from external sources.

**[0024]** The cloud **120** can be a private cloud, a public cloud, a combination of a private cloud and a public cloud (hybrid cloud), or the like. Cloud computing systems and methods abstract away physical servers, storage, networking, etc., and instead offer these as on-demand and elastic resources. The National Institute of Standards and Technology (NIST) provides a concise and specific definition which states cloud computing is a model for enabling convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. Cloud computing differs from the classic client-server model by providing applications from a server that are executed and managed by a client's web browser or the like, with no installed client version of an application required. Centralization gives cloud service providers complete control over the versions of the browser-based and other applications provided to clients, which removes the need for version upgrades or license management on individual client computing devices. The phrase "Software-as-a-Service" (SaaS) is sometimes used to describe application programs offered through cloud computing. A common shorthand for a provided cloud computing service (or even an aggregation of all existing cloud services) is "the cloud." The cloud **120** contemplates implementation via any approach known in the art.

**[0025]** The cloud **120** can be utilized to provide example cloud services, including Zscaler Internet Access (ZIA), Zscaler Private Access (ZPA), Zscaler Workload Segmentation (ZWS), and/or Zscaler Digital Experience (ZDX), all from Zscaler, Inc. (the assignee and applicant of the present application). Also, there can be multiple different clouds **120**, including ones with different architectures and multiple cloud services. The ZIA service can provide the access control, threat prevention, and data protection. ZPA can

include access control, microservice segmentation, etc. The ZDX service can provide monitoring of user experience, e.g., Quality of Experience (QoE), Quality of Service (QoS), etc., in a manner that can gain insights based on continuous, inline monitoring. For example, the ZIA service can provide a user with Internet Access, and the ZPA service can provide a user with access to enterprise resources instead of traditional Virtual Private Networks (VPNs), namely ZPA provides Zero Trust Network Access (ZTNA). Those of ordinary skill in the art will recognize various other types of cloud services are also contemplated.

### § 1.2 Zero Trust

**[0026]** FIG. 1B is a logical diagram of the cloud **120** operating as a zero-trust platform. Zero trust is a framework for securing organizations in the cloud and mobile world that asserts that no user or application should be trusted by default. Following a key zero trust principle, least-privileged access, trust is established based on context (e.g., user identity and location, the security posture of the endpoint, the app or service being requested) with policy checks at each step, via the cloud **120**. Zero trust is a cybersecurity strategy where security policy is applied based on context established through least-privileged access controls and strict user authentication—not assumed trust. A well-tuned zero trust architecture leads to simpler network infrastructure, a better user experience, and improved cyberthreat defense.

**[0027]** Establishing a zero-trust architecture requires visibility and control over the environment's users and traffic, including that which is encrypted; monitoring and verification of traffic between parts of the environment; and strong multi-factor authentication (MFA) approaches beyond passwords, such as biometrics or one-time codes. This is performed via the cloud **120**. Critically, in a zero-trust architecture, a resource's network location is not the biggest factor in its security posture anymore. Instead of rigid network segmentation, your data, workflows, services, and such are protected by software-defined micro segmentation, enabling you to keep them secure anywhere, whether in your data center or in distributed hybrid and multi-cloud environments.

**[0028]** The core concept of zero trust is simple: assume everything is hostile by default. It is a major departure from the network security model built on the centralized data center and secure network perimeter. These network architectures rely on approved IP addresses, ports, and protocols to establish access controls and validate what's trusted inside the network, generally including anybody connecting via remote access VPN. In contrast, a zero-trust approach treats all traffic, even if it is already inside the perimeter, as hostile. For example, workloads are blocked from communicating until they are validated by a set of attributes, such as a fingerprint or identity. Identity-based validation policies result in stronger security that travels with the workload wherever it communicates—in a public cloud, a hybrid environment, a container, or an on-premises network architecture.

**[0029]** Because protection is environment-agnostic, zero trust secures applications and services even if they communicate across network environments, requiring no architectural changes or policy updates. Zero trust securely connects users, devices, and applications using business policies over any network, enabling safe digital transformation. Zero trust

is about more than user identity, segmentation, and secure access. It is a strategy upon which to build a cybersecurity ecosystem.

**[0030]** At its core are three tenets:

**[0031]** Terminate every connection: Technologies like firewalls use a “passthrough” approach, inspecting files as they are delivered. If a malicious file is detected, alerts are often too late. An effective zero trust solution terminates every connection to allow an inline proxy architecture to inspect all traffic, including encrypted traffic, in real time—before it reaches its destination—to prevent ransomware, malware, and more.

**[0032]** Protect data using granular context-based policies: Zero trust policies verify access requests and rights based on context, including user identity, device, location, type of content, and the application being requested. Policies are adaptive, so user access privileges are continually reassessed as context changes.

**[0033]** Reduce risk by eliminating the attack surface: With a zero-trust approach, users connect directly to the apps and resources they need, never to networks (see ZTNA). Direct user-to-app and app-to-app connections eliminate the risk of lateral movement and prevent compromised devices from infecting other resources. Plus, users and apps are invisible to the internet, so they cannot be discovered or attacked.

### § 1.3 Log Data

**[0034]** With the cloud **120** as well as any of the network configurations **100A**, **100B**, **100C**, the log data **130** can include a rich set of statistics, logs, history, audit trails, and the like related to various endpoint **102** transactions. Generally, this rich set of data can represent activity by an endpoint **102**. This information can be for multiple endpoints **102** of a company, organization, etc., and analyzing this data can provide a wealth of information as well as training data for machine learning models.

**[0035]** The log data **130** can include a large quantity of records used in a backend data store for queries. A record can be a collection of tens of thousands of counters. A counter can be a tuple of an identifier (ID) and value. As described herein, a counter represents some monitored data associated with cybersecurity monitoring. Of note, the log data can be referred to as sparsely populated, namely a large number of counters that are sparsely populated (e.g., tens of thousands of counters or more, and possible orders of magnitude or more of which are empty). For example, a record can be stored every time period (e.g., an hour or any other time interval). There can be millions of active endpoints **102** or more. Examples of the sparsely populated log data can be the Nanolog system from Zscaler, Inc., the applicant.

**[0036]** Also, such data is described in the following:

**[0037]** Commonly-assigned U.S. Pat. No. 8,429,111, issued Apr. 23, 2013, and entitled “Encoding and compression of statistical data,” the contents of which are incorporated herein by reference, describes compression techniques for storing such logs,

**[0038]** Commonly-assigned U.S. Pat. No. 9,760,283, issued Sep. 12, 2017, and entitled “Systems and methods for a memory model for sparsely updated statistics,” the contents of which are incorporated herein by reference, describes techniques to manage sparsely updated statistics utilizing different sets of memory, hashing, memory buckets, and incremental storage, and

**[0039]** Commonly-assigned U.S. patent application Ser. No. 16/851,161, filed Apr. 17, 2020, and entitled “Systems and methods for efficiently maintaining records in a cloud-based system,” the contents of which are incorporated herein by reference, describes compression of sparsely populated log data.

**[0040]** A key aspect here is that the cybersecurity monitoring is rich and provides a wealth of information to determine various assessments of cybersecurity. In some embodiments, the log data **130** can be referred to as weblogs or the like. Of note, with various cybersecurity monitoring techniques via the network configurations **100A**, **100B**, **100C**, as well as with other network configurations, the log data **130** is a rich repository of endpoint **102** activity. Unlike websites, specific cloud services, application providers, etc., cybersecurity monitoring can log almost all of a endpoint’s **102** activity. That is, the log data **130** is not merely confined to specific activity (e.g., an endpoint’s **102** social networking activity on a specific site, an endpoint’s **102** search requests on a specific search engine, etc.), that is, when the endpoint **102** is a user.

### § 2.0 Example Server Architecture

**[0041]** FIG. 2 is a block diagram of a server **200**, which may be used as a destination on the Internet, for the network configuration **100A**, etc. The server **200** may be a digital computer that, in terms of hardware architecture, generally includes a processor **202**, input/output (I/O) interfaces **204**, a network interface **206**, a data store **208**, and memory **210**. It should be appreciated by those of ordinary skill in the art that FIG. 2 depicts the server **200** in an oversimplified manner, and a practical embodiment may include additional components and suitably configured processing logic to support known or conventional operating features that are not described in detail herein. The components (**202**, **204**, **206**, **208**, and **210**) are communicatively coupled via a local interface **212**. The local interface **212** may be, for example, but not limited to, one or more buses or other wired or wireless connections, as is known in the art. The local interface **212** may have additional elements, which are omitted for simplicity, such as controllers, buffers (caches), drivers, repeaters, and receivers, among many others, to enable communications. Further, the local interface **212** may include address, control, and/or data connections to enable appropriate communications among the aforementioned components.

**[0042]** The processor **202** is a hardware device for executing software instructions. The processor **202** may be any custom made or commercially available processor, a Central Processing Unit (CPU), an auxiliary processor among several processors associated with the server **200**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the server **200** is in operation, the processor **202** is configured to execute software stored within the memory **210**, to communicate data to and from the memory **210**, and to generally control operations of the server **200** pursuant to the software instructions. The I/O interfaces **204** may be used to receive user input from and/or for providing system output to one or more devices or components.

**[0043]** The network interface **206** may be used to enable the server **200** to communicate on a network, such as the Internet **104**. The network interface **206** may include, for

example, an Ethernet card or adapter or a Wireless Local Area Network (WLAN) card or adapter. The network interface **206** may include address, control, and/or data connections to enable appropriate communications on the network. A data store **208** may be used to store data. The data store **208** may include any volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, and the like), and combinations thereof. Moreover, the data store **208** may incorporate electronic, magnetic, optical, and/or other types of storage media. In one example, the data store **208** may be located internal to the server **200**, such as, for example, an internal hard drive connected to the local interface **212** in the server **200**. Additionally, in another embodiment, the data store **208** may be located external to the server **200** such as, for example, an external hard drive connected to the I/O interfaces **204** (e.g., SCSI or USB connection). In a further embodiment, the data store **208** may be connected to the server **200** through a network, such as, for example, a network-attached file server.

**[0044]** The memory **210** may include any volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, etc.), and combinations thereof. Moreover, the memory **210** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **210** may have a distributed architecture, where various components are situated remotely from one another but can be accessed by the processor **202**. The software in memory **210** may include one or more software programs, each of which includes an ordered listing of executable instructions for implementing logical functions. The software in the memory **210** includes a suitable Operating System (O/S) **214** and one or more programs **216**. The operating system **214** essentially controls the execution of other computer programs, such as the one or more programs **216**, and provides scheduling, input-output control, file and data management, memory management, and communication control and related services. The one or more programs **216** may be configured to implement the various processes, algorithms, methods, techniques, etc. described herein. Those skilled in the art will recognize the cloud **120** ultimately runs on one or more physical servers **200**, virtual machines, etc.

### § 3.0 Example Computing Device Architecture

**[0045]** FIG. 3 is a block diagram of a computing device **300**, which may be realized as an endpoint **102**. Specifically, the computing device **300** can form a device used by one of the endpoints **102**, and this may include common devices such as laptops, smartphones, tablets, netbooks, personal digital assistants, cell phones, e-book readers, Internet-of-Things (IoT) devices, servers, desktops, printers, televisions, streaming media devices, storage devices, and the like, i.e., anything that can communicate on a network. The computing device **300** can be a digital device that, in terms of hardware architecture, generally includes a processor **302**, I/O interfaces **304**, a network interface **306**, a data store **308**, and memory **310**. It should be appreciated by those of ordinary skill in the art that FIG. 3 depicts the computing device **300** in an oversimplified manner, and a practical embodiment may include additional components and suitably configured processing logic to support known or con-

ventional operating features that are not described in detail herein. The components (**302**, **304**, **306**, **308**, and **302**) are communicatively coupled via a local interface **312**. The local interface **312** can be, for example, but not limited to, one or more buses or other wired or wireless connections, as is known in the art. The local interface **312** can have additional elements, which are omitted for simplicity, such as controllers, buffers (caches), drivers, repeaters, and receivers, among many others, to enable communications. Further, the local interface **312** may include address, control, and/or data connections to enable appropriate communications among the aforementioned components.

**[0046]** The processor **302** is a hardware device for executing software instructions. The processor **302** can be any custom made or commercially available processor, a CPU, an auxiliary processor among several processors associated with the computing device **300**, a semiconductor-based microprocessor (in the form of a microchip or chipset), or generally any device for executing software instructions. When the computing device **300** is in operation, the processor **302** is configured to execute software stored within the memory **310**, to communicate data to and from the memory **310**, and to generally control operations of the computing device **300** pursuant to the software instructions. In an embodiment, the processor **302** may include a mobile-optimized processor such as optimized for power consumption and mobile applications. The I/O interfaces **304** can be used to receive user input from and/or for providing system output. User input can be provided via, for example, a keypad, a touch screen, a scroll ball, a scroll bar, buttons, a barcode scanner, and the like. System output can be provided via a display device such as a Liquid Crystal Display (LCD), touch screen, and the like.

**[0047]** The network interface **306** enables wireless communication to an external access device or network. Any number of suitable wireless data communication protocols, techniques, or methodologies can be supported by the network interface **306**, including any protocols for wireless communication. The data store **308** may be used to store data. The data store **308** may include any volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, and the like)), nonvolatile memory elements (e.g., ROM, hard drive, tape, CDROM, and the like), and combinations thereof. Moreover, the data store **308** may incorporate electronic, magnetic, optical, and/or other types of storage media.

**[0048]** The memory **310** may include any volatile memory elements (e.g., random access memory (RAM, such as DRAM, SRAM, SDRAM, etc.)), nonvolatile memory elements (e.g., ROM, hard drive, etc.), and combinations thereof. Moreover, the memory **310** may incorporate electronic, magnetic, optical, and/or other types of storage media. Note that the memory **310** may have a distributed architecture, where various components are situated remotely from one another, but can be accessed by the processor **302**. The software in memory **310** can include one or more software programs, each of which includes an ordered listing of executable instructions for implementing logical functions. In the example of FIG. 3, the software in the memory **310** includes a suitable operating system **314** and programs **316**. The operating system **314** essentially controls the execution of other computer programs and provides scheduling, input-output control, file and data management, memory management, and communication

control and related services. The programs **316** may include various applications, add-ons, etc. configured to provide end-user functionality with the computing device **300**. For example, example programs **316** may include, but not limited to, a web browser, social networking applications, streaming media applications, games, mapping and location applications, electronic mail applications, financial applications, and the like. The connector application **110** can be one of the example programs.

#### § 4.0 Application for Traffic Forwarding and Monitoring

[0049] Again, the network configuration **100B** includes a connector application **110** that is executed on the computing device **300**. The connector application **110** can perform similar functionality as the server **200**, as well as coordinated functionality with the server **200** (a combination of the network configurations **100A**, **100B**). Of course, various embodiments are contemplated herein, including combinations of the network configurations **100A**, **100B**, **100C** together. For example, the connector application **110** can perform similar functionality as the cloud **120**, as well as coordinated functionality with the cloud **120**.

[0050] FIG. 4 is a network diagram of an exemplary network configuration illustrating a connector application **110** on computing devices **300** configured to operate through the cloud **120**. Different types of computing devices **300** are proliferating, including Bring Your Own Device (BYOD) as well as IT-managed devices. The conventional approach for a computing device **300** to operate with the cloud **120** as well as for accessing enterprise resources includes complex policies, VPNs, poor user experience, etc. The connector application **110** can automatically forward user traffic with the cloud **120** as well as ensuring that security and access policies are enforced, regardless of device, location, operating system, or application. The connector application **110** automatically determines if an endpoint **102** is looking to access the open Internet **104**, a SaaS app, or an internal app running in public, private, or the datacenter and routes mobile traffic through the cloud **120**. The connector application **110** can support various cloud services, including ZIA, ZPA, ZDX, etc., allowing the best in class security with zero trust access to internal applications. As described herein, the connector application **110** can also be referred to as a connector application.

[0051] The connector application **110** is configured to auto-route traffic for seamless user experience. This can be protocol as well as application-specific, and the connector application **110** can route traffic with a nearest or best fit node of the cloud **120**. Further, the connector application **110** can detect trusted networks, allowed applications, etc. and support secure network access. The connector application **110** can also support the enrollment of the computing device **300** prior to accessing applications, the internet, or any services provided by the cloud **120**. The connector application **110** can uniquely detect the endpoints **102** based on fingerprinting the computing device **300**, using criteria like device model, platform, operating system, device posture, etc. The connector application **110** can support Mobile Device Management (MDM) functions, allowing IT personnel to deploy and manage the computing devices **300** seamlessly. This can also include the automatic installation of client and SSL certificates during enrollment. Finally, the connector application **110** provides visibility into device and

app usage of the computing device **300**, when a user is the endpoint **102** utilizing the computing device **300**.

[0052] The connector application **110** supports a secure, lightweight tunnel between the computing device **300** and the cloud **120**. For example, the lightweight tunnel can be HTTP-based. With the connector application **110**, there is no requirement for PAC files, an IPsec VPN, authentication cookies, or user setup.

#### § 5.0 Endpoint Process Based Policy Enforcement and Logging

[0053] As described herein, the cloud **120** is adapted to monitor, log, and enforce policy on traffic inline. As part of offering cybersecurity through the example network configurations **100A**, **100B**, **100C**, a large amount of cybersecurity data is obtained. This cybersecurity data can include user data, device data, traffic data (such as user-destination transactions), and the like. Various embodiments of the present disclosure focus on using this cybersecurity data along with a customer's data to perform various security tasks including developing customer machine learning models and enforcing policies.

[0054] FIG. 5 is a flow diagram of a network configuration including an endpoint **102**, a connector application **110**, and the cloud **120**. Again, the endpoint **102** can be a user utilizing a computing device **300**, an Internet-of-Things (IoT) device, and any customer specific endpoint such as infrastructure and the like adapted to communicate through the cloud **120**. Again, an endpoint **102** can be configured to have the connector application **110** executing thereon. The connector application **110** is adapted to act, and can be contemplated, as a gateway for the endpoint for forwarding traffic to the cloud **120**, enforce policy, collect endpoint **102** data, and the like. Logically, the cloud **120** can be viewed as an overlay network between one or more endpoints **102** and one or more destinations such as the Internet **104**, cloud services, SaaS, etc. as described the cloud **120** is adapted to monitor traffic between the one or more endpoints and destinations inline.

[0055] Endpoints **102** can have various Endpoint Processes (EP) **140** operating thereon. For example, end endpoint **102** being a user device can have various applications installed thereon which can be utilized to communicate with the internet **104**. In this example, the various applications are contemplated as the endpoint processes **140**. In various embodiments, an endpoint process **140** can be any software installed on an endpoint **102** which is adapted to be used to access the internet **104**, cloud **120**, etc. For example, the present disclosure contemplates an endpoint process **140** as an application, a web browser, or any other process or software of the like. Thus, the connector application **110** is adapted to intercept traffic from any endpoint process **140** and forward the traffic to the cloud **120**, or any of the cloud-provided cybersecurity services described herein, in-band with the traffic.

[0056] The present disclosure proposes various systems and methods for extracting additional information from the various endpoint processes **140** found on an endpoint **102**. For example, when a new traffic flow is initiated from the endpoint **102** to the cloud **120**, via the connector application **110**, the connector application **110** can be adapted to also send endpoint process **140** specific metadata along with the traffic. In one use case, when the endpoint **102** is a user device operating a web browser, responsive to a new flow of

traffic being initiated from the web browser, the connector application 110 can extract metadata from the web browser and forward the extracted web browser metadata to the cloud 120. In this example, the user device is the endpoint 102 and the web browser is the endpoint process 140. The cloud 120 can identify endpoint process metadata within the traffic and perform the processing of the metadata based thereon. The cloud 120 can then process the metadata and enforce policy, and/or log the data for future use. Based on the processing of the metadata, the cloud 120 and/or the connector application 110 can be adapted to perform one or more actions. These actions include allowing the traffic, blocking the traffic, alerting a user of the endpoint 102, etc. The alerting can include instructing the user to update software of the endpoint process, use a different endpoint process, etc. (further described herein). In various embodiments, the alerting can be provided via the application 110 executing on the user device, for example via a push notification or the like. Additionally, an alert can include instructions for remediating the alert, again, for example, providing instructions for how to update the software of an endpoint process, or providing recommendations for other endpoint processes available on the endpoint.

**[0057]** In various embodiments, the cloud 120 is adapted to log the metadata in log data 130. The logged metadata can then be used by the cloud 120 to train various Artificial Intelligence (AI) models, Machine Learning (ML) models etc. The logging and utilization of the metadata can be performed on a per-customer basis. That is, metadata can be logged and kept separate for each customer of the cloud 120 and its various security services. Similarly, the cloud 120 can enforce metadata-based policy (further described herein) on a per-customer basis. For example, different customers of the cloud 120 can have different policies enforced in their respective environments. Based thereon, the cloud 120 can enforce these policies based on the metadata and the customer with which the endpoint 102 is associated with.

**[0058]** In various embodiments, the metadata collected by the connector application 110 from the endpoint process 140 can include one or more of signer name, signature timestamp, product name, product version, original file name, executable hash, Process ID (PID), process name, process owner, process running duration, and the like. The signer name can be a company which authored the piece of software (endpoint process 140), for example, Google would be the signer name of the Chrome web browser application. A signature timestamp can certify a processes execution at a point in time with added levels of assurance. The product name can be the name of the endpoint process being used for the request. A product version can also be extracted and forwarded to the cloud 120 for determination that a latest version or an approved version of the endpoint process 140 is being utilized. The original file name can be utilized to determine the name issued to the process when it was created, thus can be utilized to determine that the endpoint process 140 has not been renamed. Further, the executable hash or a cryptographic signature can be utilized to determine if the endpoint process 140 matches what is expected and verify that the endpoint process 140 is legitimate. Additionally, the PID, process name, process owner, and other properties of an endpoint process 140 can be forwarded to the cloud 120 for identification purposes and the like. The process running duration can be utilized to deter-

mine how long an endpoint process has been running on the endpoint 102, which can be utilized in policy enforcement and risk assessments.

**[0059]** By utilizing the present systems and methods, zero trust policy can be extended down to an application level. That is, based on the metadata forwarded by the connector application 110, the cloud 120 can determine what endpoint process 140 an endpoint 102 is using, and can thus enforce zero trust policy based thereon. For example, various endpoint processes 140 can be trusted whereas others can be untrusted, and traffic can be either allowed or blocked based thereon. In a use case, if an endpoint is attempting to access a sensitive destination, i.e., a destination having sensitive data therein, the cloud 120 can restrict access to such destinations based on the endpoint process 140 which the endpoint 102 is using. Similarly, the cloud 120 can restrict access and/or only allow access to various destinations based on the endpoint process 140 identified via the endpoint process metadata.

**[0060]** For example, the cloud 120 can be configured to only allow access to sensitive destinations via a trusted corporate provided web browser of a latest or approved version regardless of header information. In this example, the cloud 120 processes the endpoint process metadata sent in-band with the traffic/request and determines that the endpoint process 140 used for the request is the corporate provided web browser, and is the latest or approved version of the corporate provided web browser. The cloud 120 can then allow the request/access based thereon. Similarly, the cloud 120 can block a request if it finds that either the endpoint process used to access the destination is not the trusted corporate provided web browser, or if the trusted corporate provided web browser is not of the latest version or an approved version. In such cases, the cloud 120 can communicate a reason for blocking the request to the connector application 110, which can then provide a UI for a user of the endpoint explaining the reason for blocking the request. Users can then solve any endpoint process deficiencies by using a different endpoint process, updating the endpoint process, etc. Additionally, the connector application 110 can be adapted to automatically update an endpoint process such as a corporate provided web browser responsive to receiving a reason of the blocking from the cloud 120, thereby resolving the issue without user intervention. This process is useful for endpoints which users do not have access to on a regular basis.

**[0061]** In various embodiments, the present systems can utilize the provided endpoint process metadata for determining if Transport Layer Security (TLS) decryption, or other inspection feature, is necessary. For example, based on the endpoint process metadata of the endpoint process originating the request, the present systems can choose to bypass TLS decryption or bypass any other cloud inspection engines. In one use case, if the systems determine that the endpoint process originating a request is a corporate provided web browser of the latest or approved version and does not show any signs of maliciousness, i.e., based on the executable hash or cryptographic signature, the systems can decide to bypass TLS decryption or any other cloud inspection engines. Again, the risk associated with the destination of the request can also weigh into such decisions. For example, the destination can be considered a sensitive destination or not a sensitive destination, the systems can use such information in the deciding.

**[0062]** By utilizing the endpoint process metadata, more granular data protection can be enforced by the cloud **120**. That is, the cloud **120** can enforce specific policies and/or services such as DLP, access control policies, etc. based on the endpoint process **140** which the endpoint **102** is using for the connection. For example, if an endpoint is using an unsanctioned application for accessing the internet, the cloud **120** can identify this, and enforce more strict DLP policy for the traffic. This provides an additional dimension for data protection and security.

**[0063]** In addition to utilizing the metadata for enforcing policy on traffic, the present systems and methods include utilizing the metadata for visibility and inventory management. Many organizations have difficulty managing their software, endpoints, resources, etc. This is because they do not have good visibility that would stop them from having to inventory everything they have. Alternatively, it is much more useful to inventory everything that is in use, i.e., endpoint processes **140** in use by their endpoints **102**. That is, the present systems can utilize the metadata forwarded to the cloud **120** to log and inventory all applications, software, etc. running on, and sending traffic from a customer's endpoints. The cloud **120** can then provide each customer an inventory of all endpoint processes **140** being utilized on their endpoints **102** based on the logging.

**[0064]** Again, the cloud **120** can enforce user application and web browser version control. Many vulnerabilities can be introduced because of users not updating their software on time. Thus, the present systems introduce another level of protection by enforcing software version control. For example, the cloud **120** can block an endpoint from browsing the internet unless the web browser, or other software, used for the connection is updated.

**[0065]** For endpoints **102** such as IoT devices, because such endpoints **102** run very specific standard software, the systems can utilize the metadata for verification that the software is up to date, the correct software, not tampered with, etc.

#### § 5.1 Process for Endpoint Process Based Policy Enforcement and Logging

**[0066]** FIG. 6 is a flow chart of a process **600** for endpoint application metadata based policy enforcement. The process **600** includes monitoring traffic via a cloud, the traffic being monitored inline between one or more endpoints and one or more destinations (step **602**); identifying, within a request from an endpoint, endpoint process metadata associated with an endpoint process used to make the request (step **604**); processing the endpoint process metadata (step **606**); and performing one or more actions on the request based on the processing (step **608**).

**[0067]** The process **600** can further include wherein the one or more actions include blocking the request, allowing the request, and alerting a user of the endpoint. The request can be forwarded to the cloud via a connector application executing on the endpoint. The connector application can be adapted to collect endpoint process metadata from an endpoint process used to make the request. The connector application can be adapted to forward the endpoint process metadata to the cloud in-band with the request. The endpoint process metadata can include one or more of a signor name, signature timestamp, product name, product version, original file name, executable hash, Process Identity (PID), process name, process owner, process running duration, and

other properties of the endpoint process. When the endpoint process metadata includes the product version of the endpoint process, the steps can further include processing the endpoint process metadata to determine if the product version of the endpoint process is the latest version available or an approved version; and performing one or more actions on the request based thereon. Responsive to determining the product version of the endpoint process is not the latest version available or an approved version, the steps can include blocking the request and alerting a user of the endpoint. The steps can further include logging the endpoint process metadata. The steps can further include providing a customer of the cloud with an inventory of all endpoint processes being utilized on their endpoints based on the logging.

#### \$6.0 Conclusion

**[0068]** It will be appreciated that some embodiments described herein may include one or more generic or specialized processors (“one or more processors”) such as microprocessors; Central Processing Units (CPUs); Digital Signal Processors (DSPs); customized processors such as Network Processors (NPs) or Network Processing Units (NPU), Graphics Processing Units (GPUs), or the like; Field Programmable Gate Arrays (FPGAs); and the like along with unique stored program instructions (including software and/or firmware) for control thereof to implement, in conjunction with certain non-processor circuits, some, most, or all of the functions of the methods and/or systems described herein. Alternatively, some or all functions may be implemented by a state machine that has no stored program instructions, or in one or more Application-Specific Integrated Circuits (ASICs), in which each function or some combinations of certain of the functions are implemented as custom logic or circuitry. Of course, a combination of the aforementioned approaches may be used. For some of the embodiments described herein, a corresponding device in hardware and optionally with software, firmware, and a combination thereof can be referred to as “circuitry configured or adapted to,” “logic configured or adapted to,” “a circuit configured to,” “one or more circuits configured to,” etc. perform a set of operations, steps, methods, processes, algorithms, functions, techniques, etc. on digital and/or analog signals as described herein for the various embodiments.

**[0069]** Moreover, some embodiments may include a non-transitory computer-readable storage medium having computer-readable code stored thereon for programming a computer, server, appliance, device, processor, circuit, etc. each of which may include a processor to perform functions as described and claimed herein. Examples of such computer-readable storage mediums include, but are not limited to, a hard disk, an optical storage device, a magnetic storage device, a Read-Only Memory (ROM), a Programmable Read-Only Memory (PROM), an Erasable Programmable Read-Only Memory (EPROM), an Electrically Erasable Programmable Read-Only Memory (EEPROM), Flash memory, and the like. When stored in the non-transitory computer-readable medium, software can include instructions executable by a processor or device (e.g., any type of programmable circuitry or logic) that, in response to such execution, cause a processor or the device to perform a set

of operations, steps, methods, processes, algorithms, functions, techniques, etc. as described herein for the various embodiments.

**[0070]** Although the present disclosure has been illustrated and described herein with reference to embodiments and specific examples thereof, it will be readily apparent to those of ordinary skill in the art that other embodiments and examples may perform similar functions and/or achieve like results. All such equivalent embodiments and examples are within the spirit and scope of the present disclosure, are contemplated thereby, and are intended to be covered by the following claims. Further, the various elements, operations, steps, methods, processes, algorithms, functions, techniques, modules, circuits, etc. described herein contemplate use in any and all combinations with one another, including individually as well as combinations of less than all of the various elements, operations, steps, methods, processes, algorithms, functions, techniques, modules, circuits, etc.

What is claimed is:

1. A method comprising steps of:  
monitoring traffic via a cloud, the traffic being monitored inline between one or more endpoints and one or more destinations;  
identifying, within a request from an endpoint, endpoint process metadata associated with an endpoint process used to make the request;  
processing the endpoint process metadata; and  
performing one or more actions on the request based on the processing.
2. The method of claim 1, wherein the one or more actions include blocking the request, allowing the request, bypassing Transport Layer Security (TLS) decryption for the request, and alerting a user of the endpoint.
3. The method of claim 1, wherein the request is forwarded to the cloud via a connector application executing on the endpoint.
4. The method of claim 3, wherein the connector application is adapted to collect endpoint process metadata from an endpoint process used to make the request.
5. The method of claim 4, wherein the connector application is adapted to forward the endpoint process metadata to the cloud in-band with the request.
6. The method of claim 1, wherein the endpoint process metadata includes one or more of a signor name, signature timestamp, product name, product version, original file name, executable hash, Process Identity (PID), process name, process owner, and process running duration of the endpoint process.
7. The method of claim 6, wherein the endpoint process metadata includes the product version of the endpoint process, and wherein the steps comprise:  
processing the endpoint process metadata to determine if the product version of the endpoint process is a latest version available; and  
performing one or more actions on the request based thereon.
8. The method of claim 7, wherein responsive to determining the product version of the endpoint process is not the latest version available, blocking the request and alerting a user of the endpoint.
9. The method of claim 1, wherein the steps further comprise:  
logging the endpoint process metadata.

10. The method of claim 9, wherein the steps further comprise:

providing a customer of the cloud with an inventory of all endpoint processes being utilized on their endpoints based on the logging.

11. A non-transitory computer-readable medium comprising instructions that, when executed, cause one or more processors to perform steps of:

monitoring traffic via a cloud, the traffic being monitored inline between one or more endpoints and one or more destinations;

identifying, within a request from an endpoint, endpoint process metadata associated with an endpoint process used to make the request;

processing the endpoint process metadata; and  
performing one or more actions on the request based on the processing.

12. The non-transitory computer-readable medium of claim 11, wherein the one or more actions include blocking the request, allowing the request, bypassing Transport Layer Security (TLS) decryption for the request, and alerting a user of the endpoint.

13. The non-transitory computer-readable medium of claim 11, wherein the request is forwarded to the cloud via a connector application executing on the endpoint.

14. The non-transitory computer-readable medium of claim 13, wherein the connector application is adapted to collect endpoint process metadata from an endpoint process used to make the request.

15. The non-transitory computer-readable medium of claim 14, wherein the connector application is adapted to forward the endpoint process metadata to the cloud in-band with the request.

16. The non-transitory computer-readable medium of claim 11, wherein the endpoint process metadata includes one or more of a signor name, signature timestamp, product name, product version, original file name, executable hash, Process Identity (PID), process name, process owner, and process running duration of the endpoint process.

17. The non-transitory computer-readable medium of claim 16, wherein the endpoint process metadata includes the product version of the endpoint process, and wherein the steps comprise:

processing the endpoint process metadata to determine if the product version of the endpoint process is a latest version available; and

performing one or more actions on the request based thereon.

18. The non-transitory computer-readable medium of claim 17, wherein responsive to determining the product version of the endpoint process is not the latest version available, blocking the request and alerting a user of the endpoint.

19. The non-transitory computer-readable medium of claim 11, wherein the steps further comprise:

logging the endpoint process metadata.

20. The non-transitory computer-readable medium of claim 19, wherein the steps further comprise:

providing a customer of the cloud with an inventory of all endpoint processes being utilized on their endpoints based on the logging.

\* \* \* \* \*