



US 20250260752A1

(19) **United States**

(12) **Patent Application Publication**
Smith et al.

(10) **Pub. No.: US 2025/0260752 A1**

(43) **Pub. Date: Aug. 14, 2025**

(54) **TRANSACTION PROCESSING SYSTEM
WITH DECOMPOSED ARCHITECTURE**

(52) **U.S. Cl.**

CPC **H04L 67/63** (2022.05); **H04L 67/306**
(2013.01)

(71) Applicant: **Games Global Operations Limited,**
Douglas (IM)

(57)

ABSTRACT

(72) Inventors: **Grant Michael Smith,** Union Mills
(IM); **Scott Xavier Dukes,**
KwaZulu-Natal (ZA); **Matthew John
Harris,** Berkhamsted (GB)

An implementation may involve: receiving, by an application launcher, a message on behalf of a client device; determining a server application based on the message and metadata, wherein the server application is associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions, wherein a generate result step causes the server application to determine an outcome of executing the path, and wherein a handle request step causes the server application to provide, to an external service, a representation of a request related to the outcome; providing, by the application launcher, a client application to the client device; and executing, by an orchestrator, the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.

(21) Appl. No.: **19/190,031**

(22) Filed: **Apr. 25, 2025**

Related U.S. Application Data

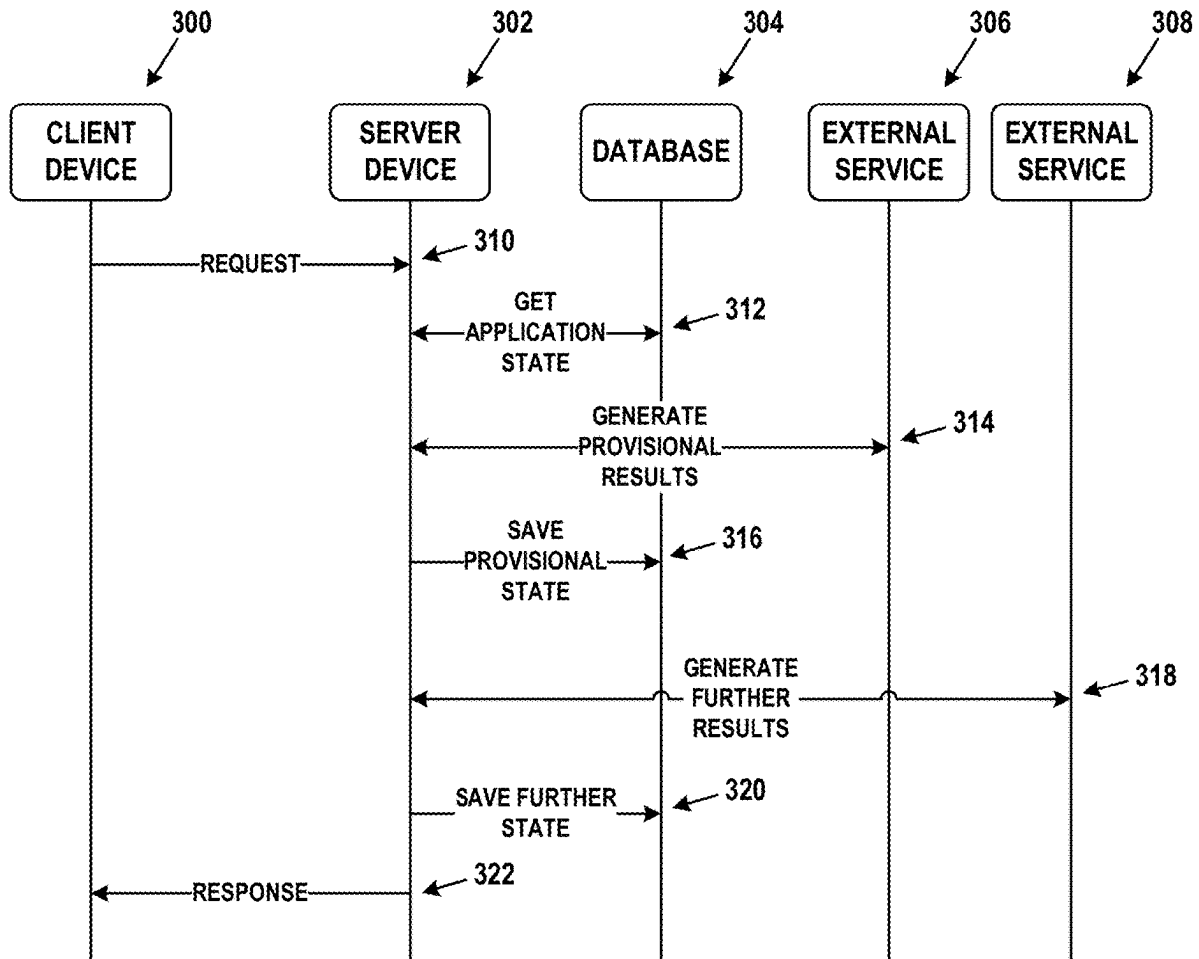
(63) Continuation of application No. 18/523,161, filed on
Nov. 29, 2023, now Pat. No. 12,316,720.

Publication Classification

(51) **Int. Cl.**

H04L 67/63 (2022.01)

H04L 67/306 (2022.01)






FIG. 1




FIG. 2

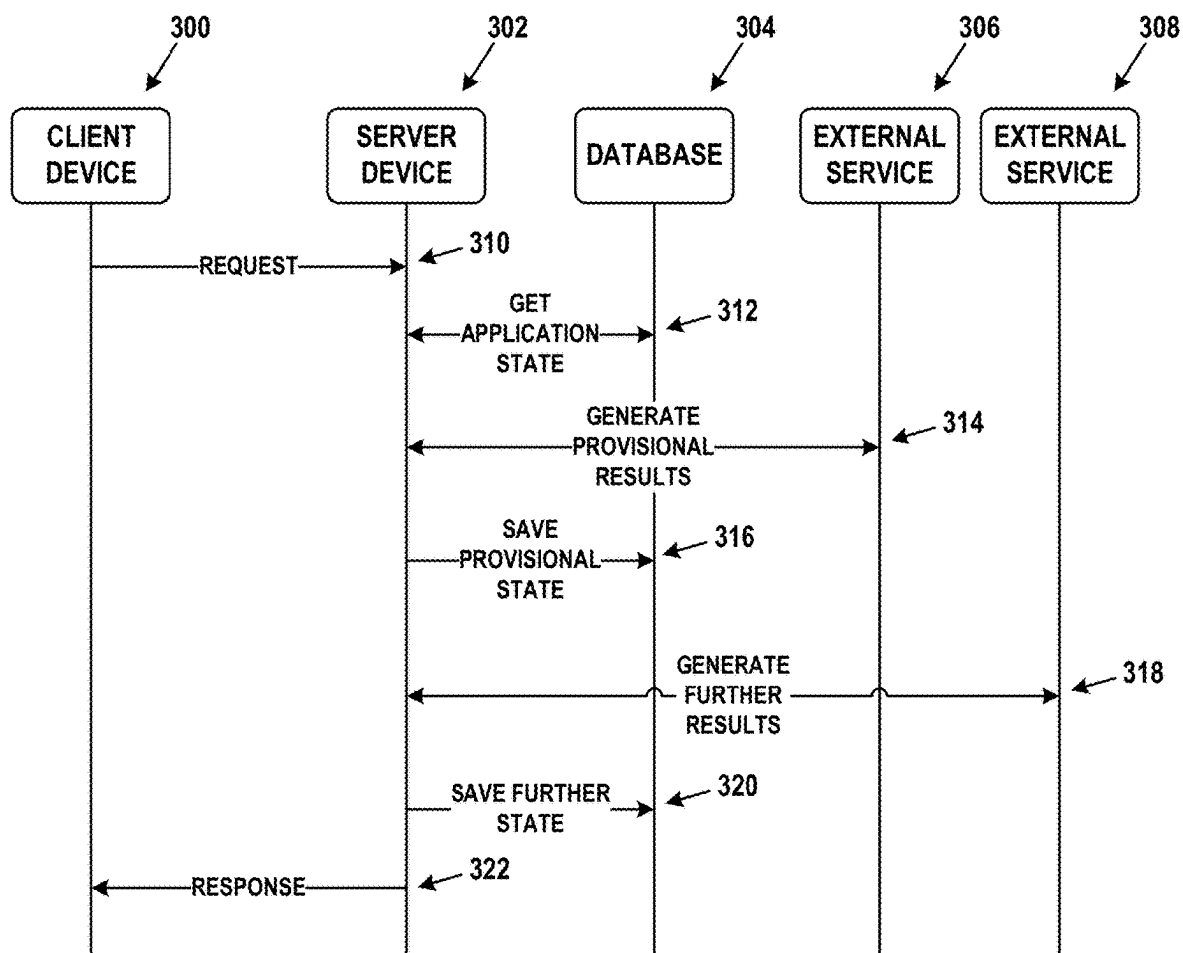
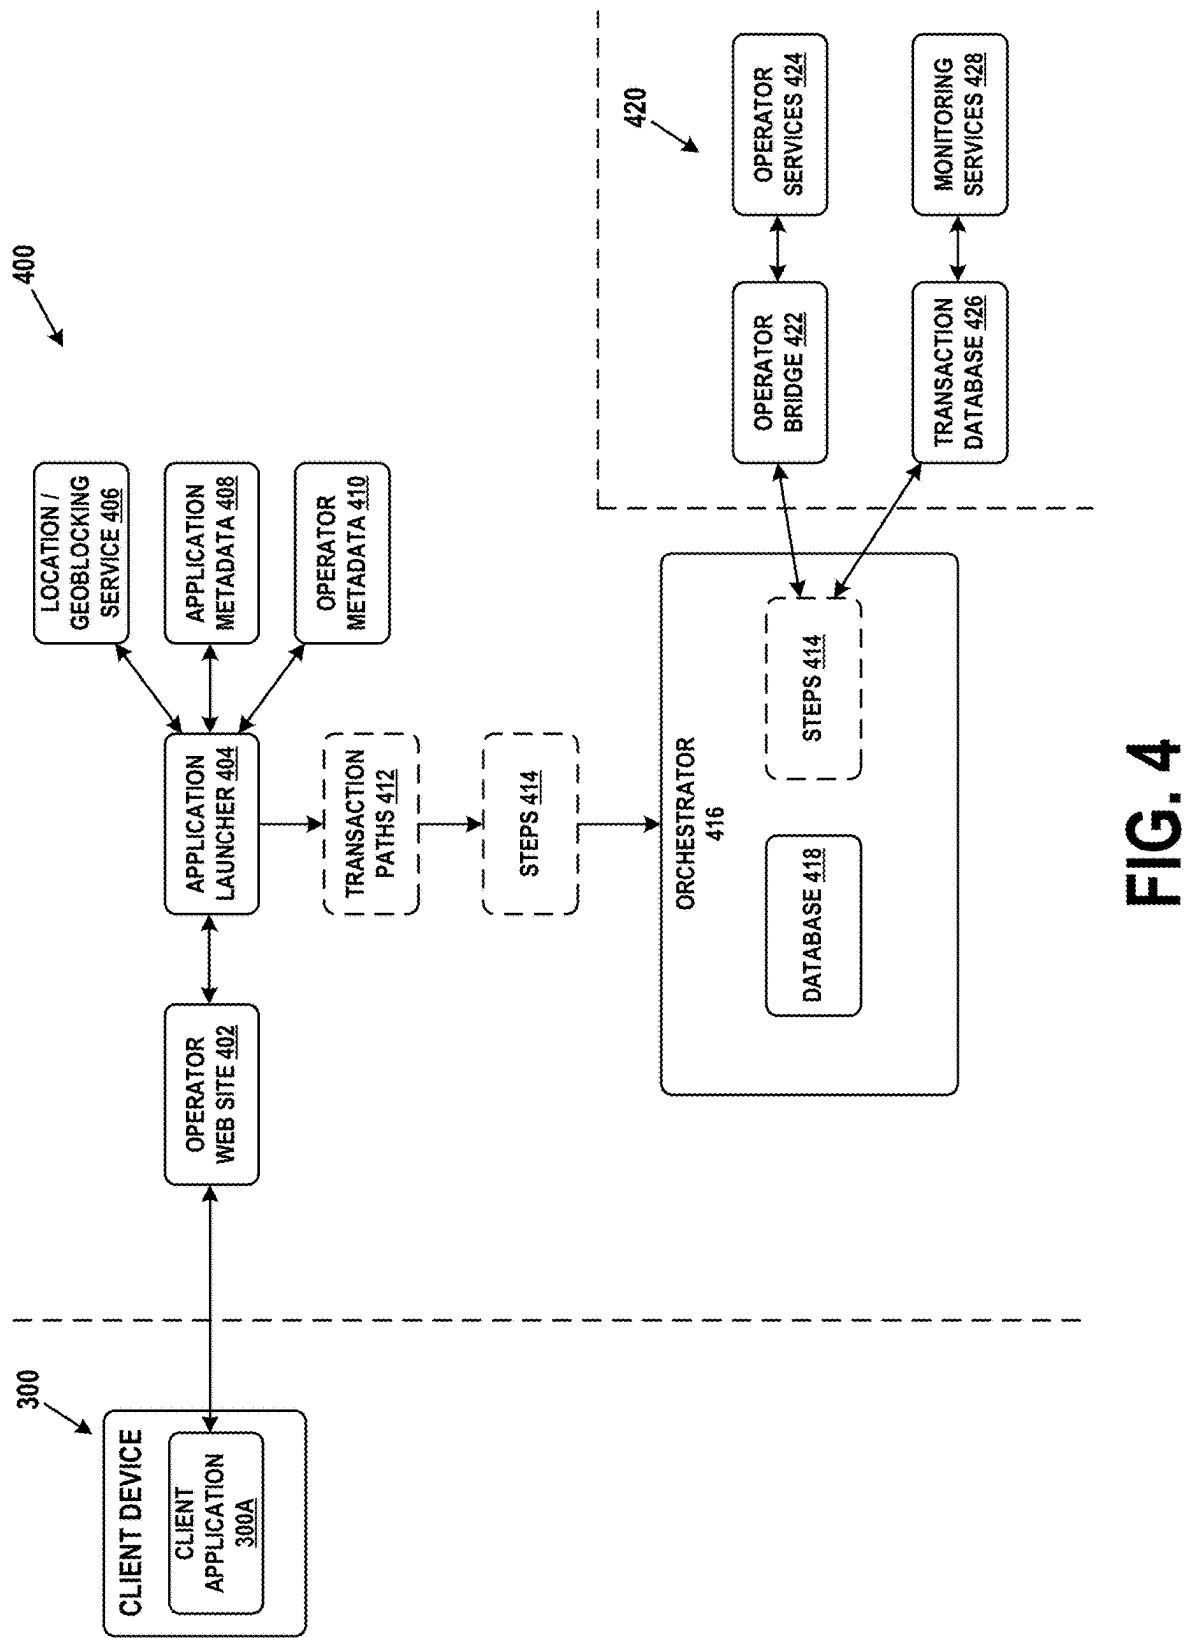


FIG. 3



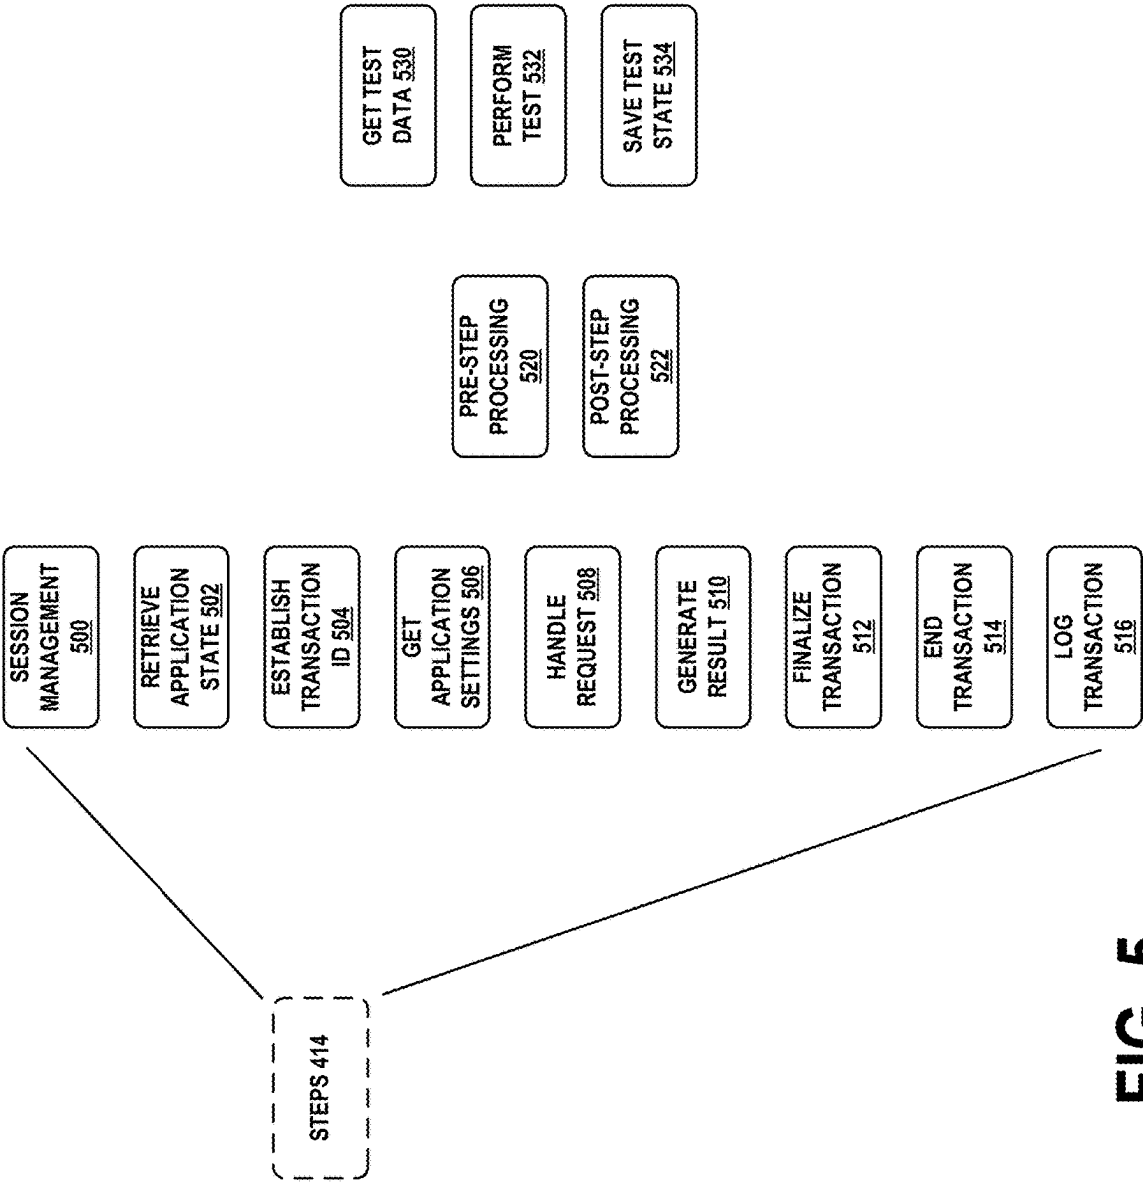


FIG. 5




FIG. 6




FIG. 7

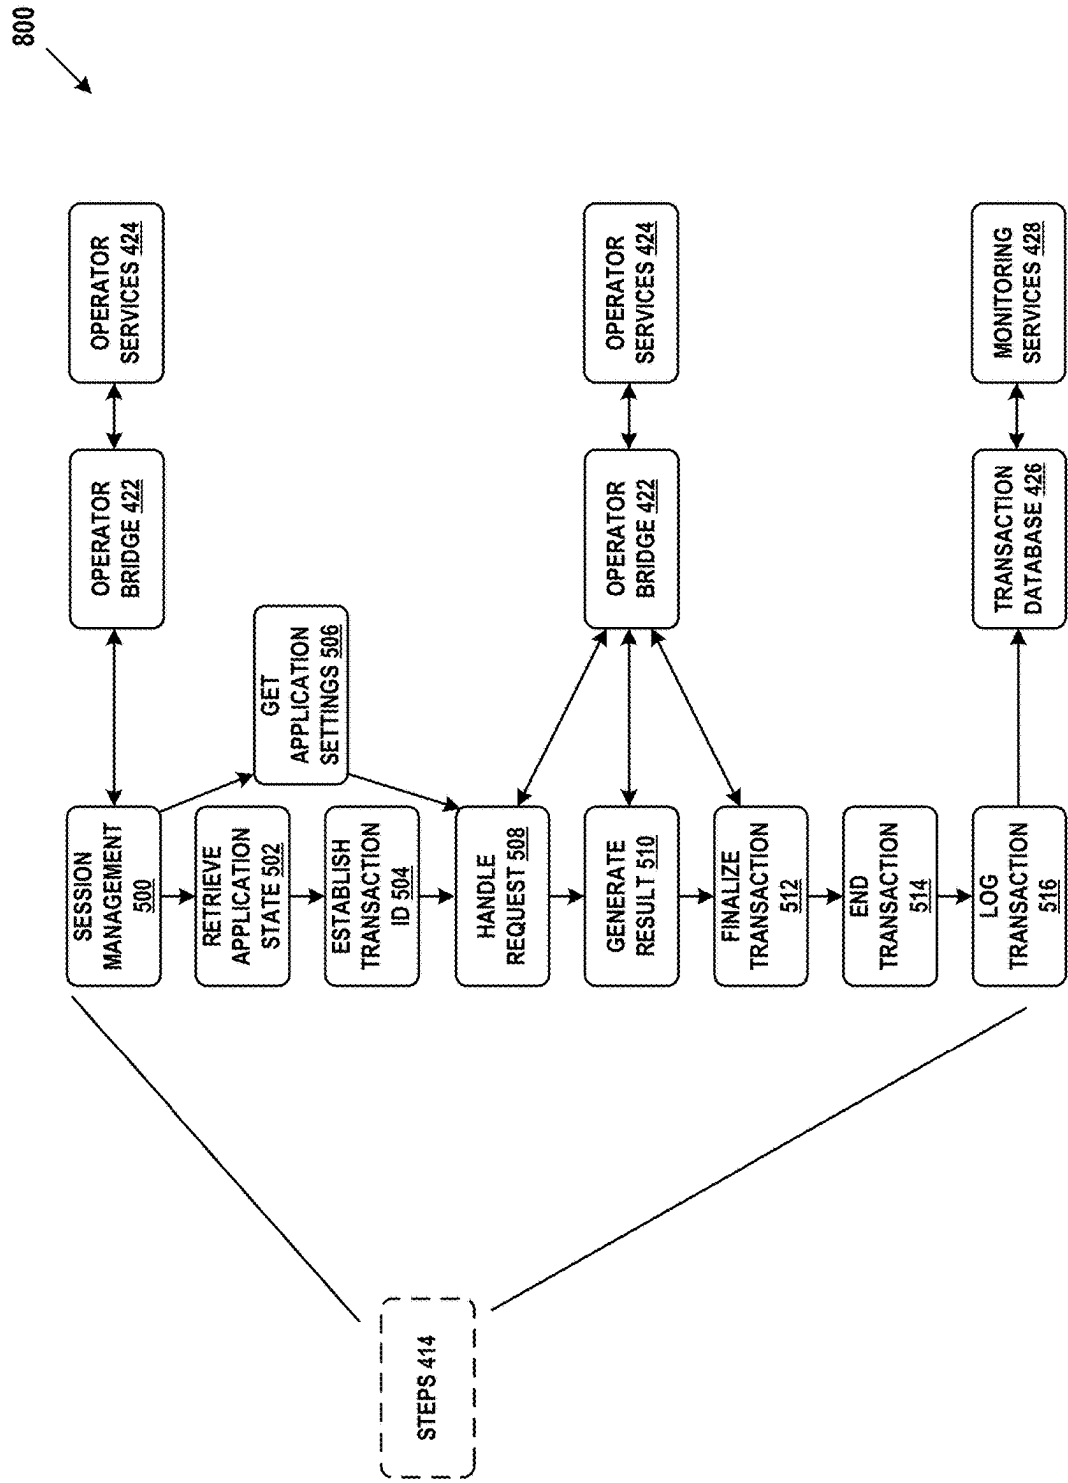


FIG. 8

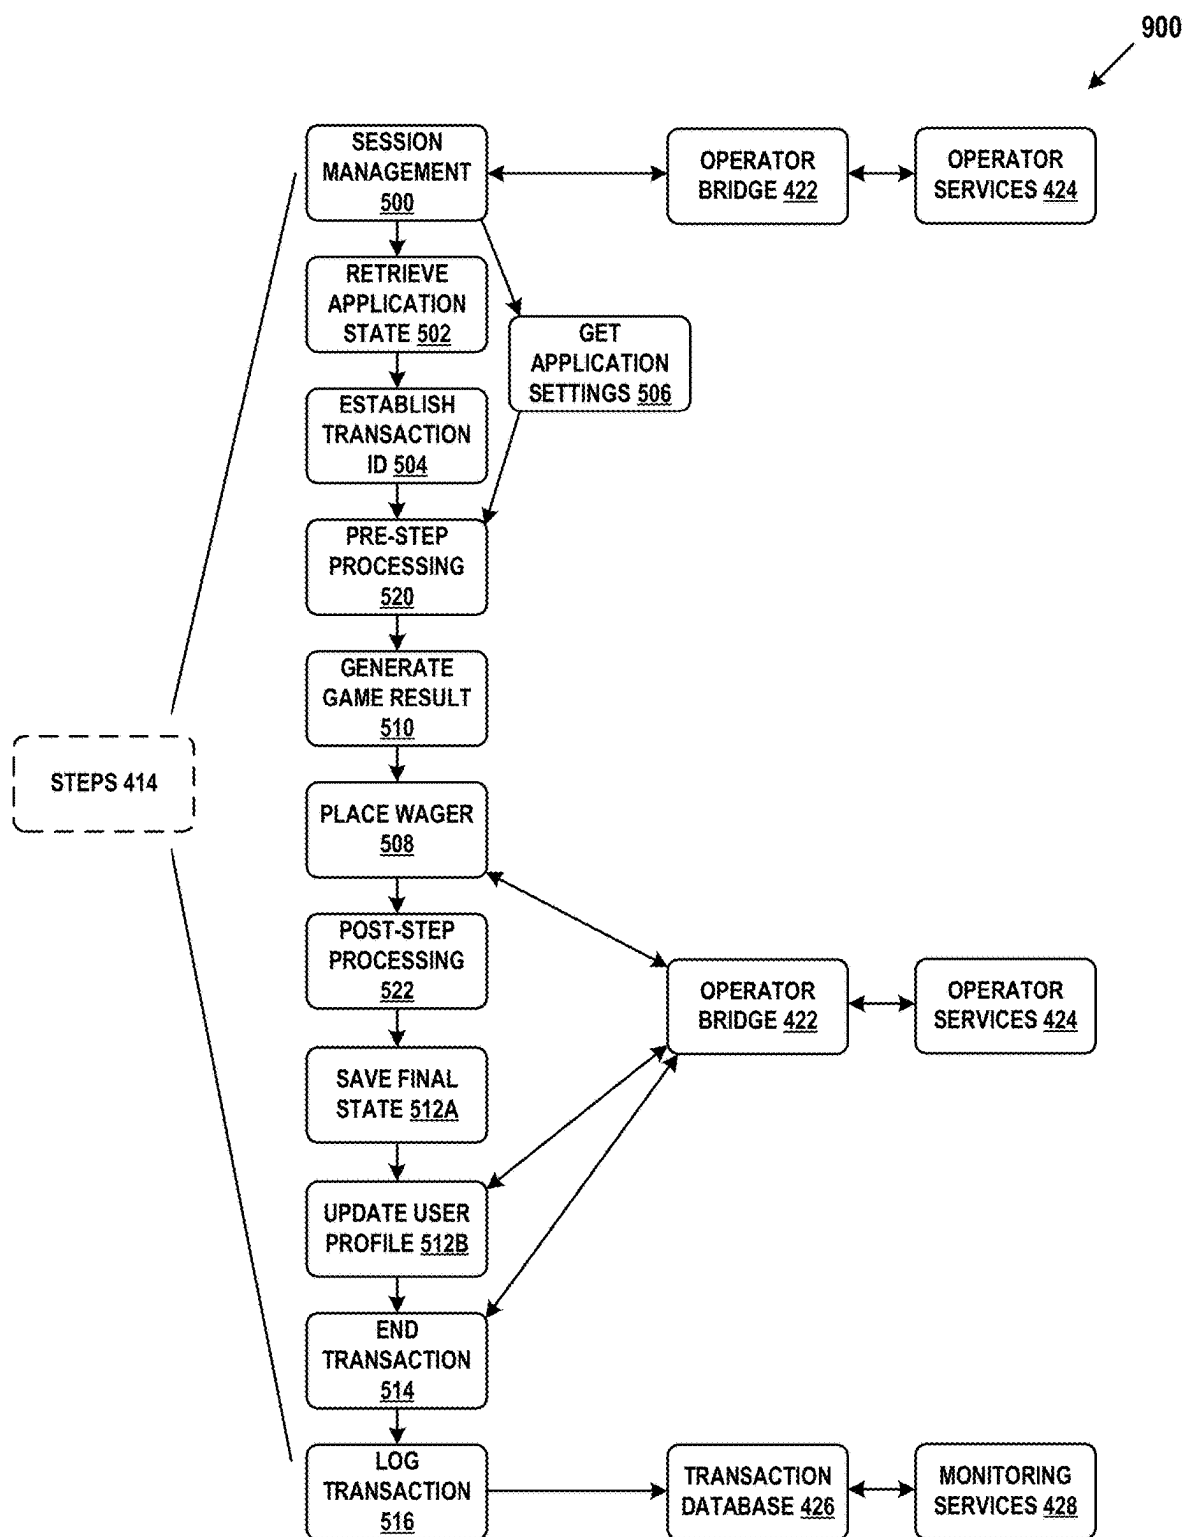


FIG. 9

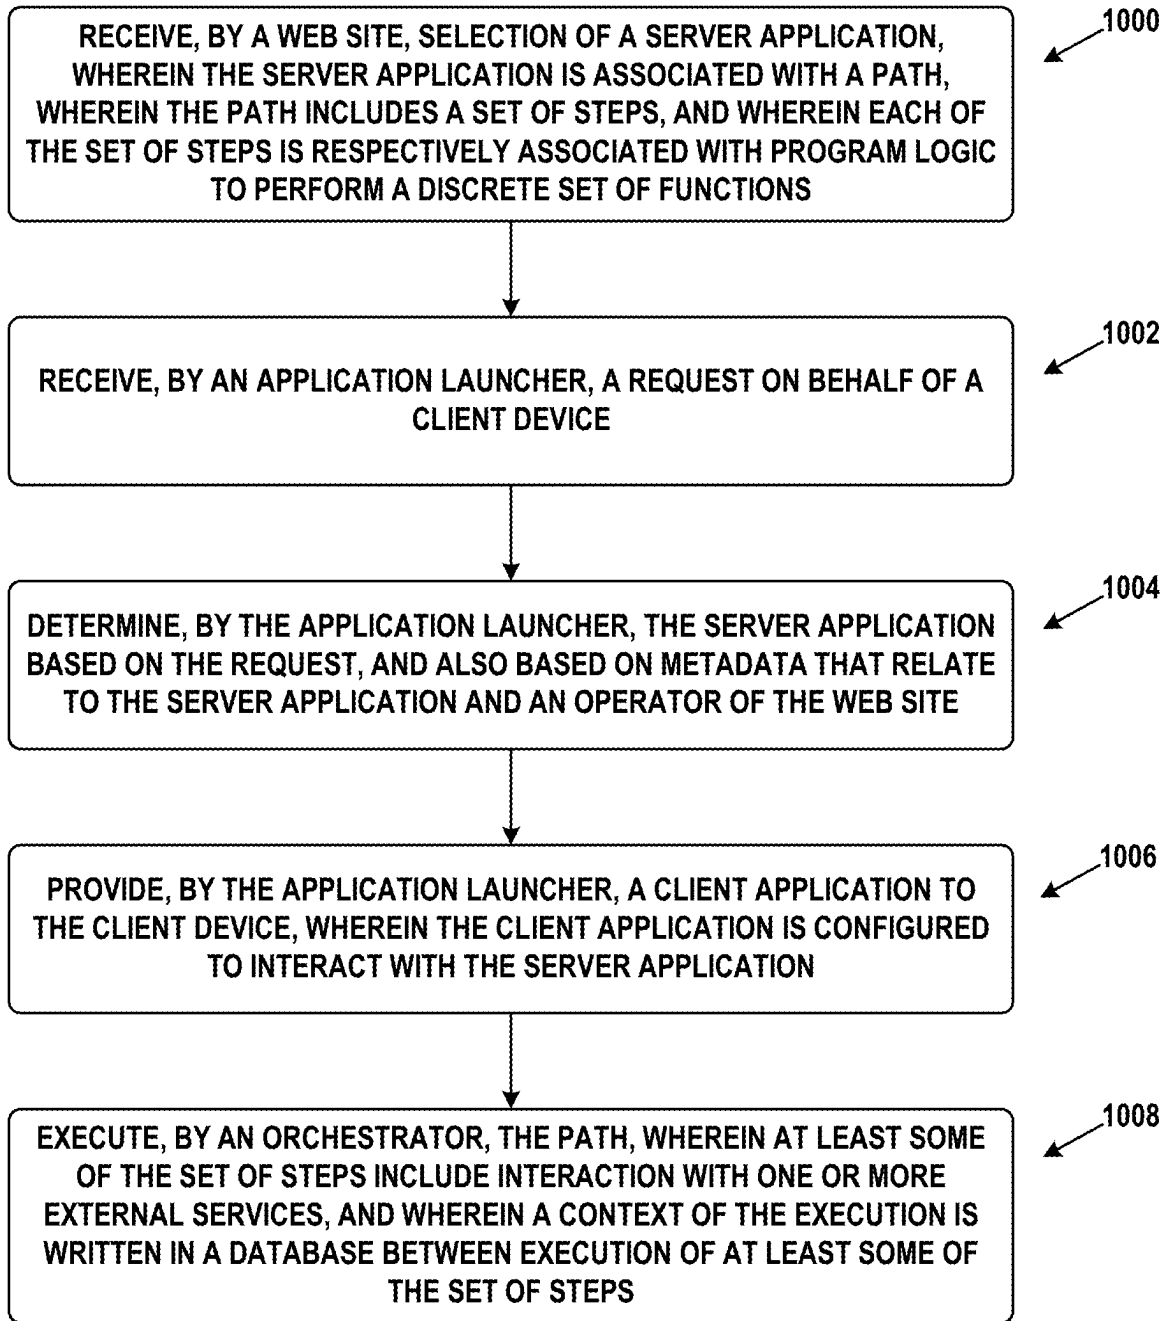


FIG. 10

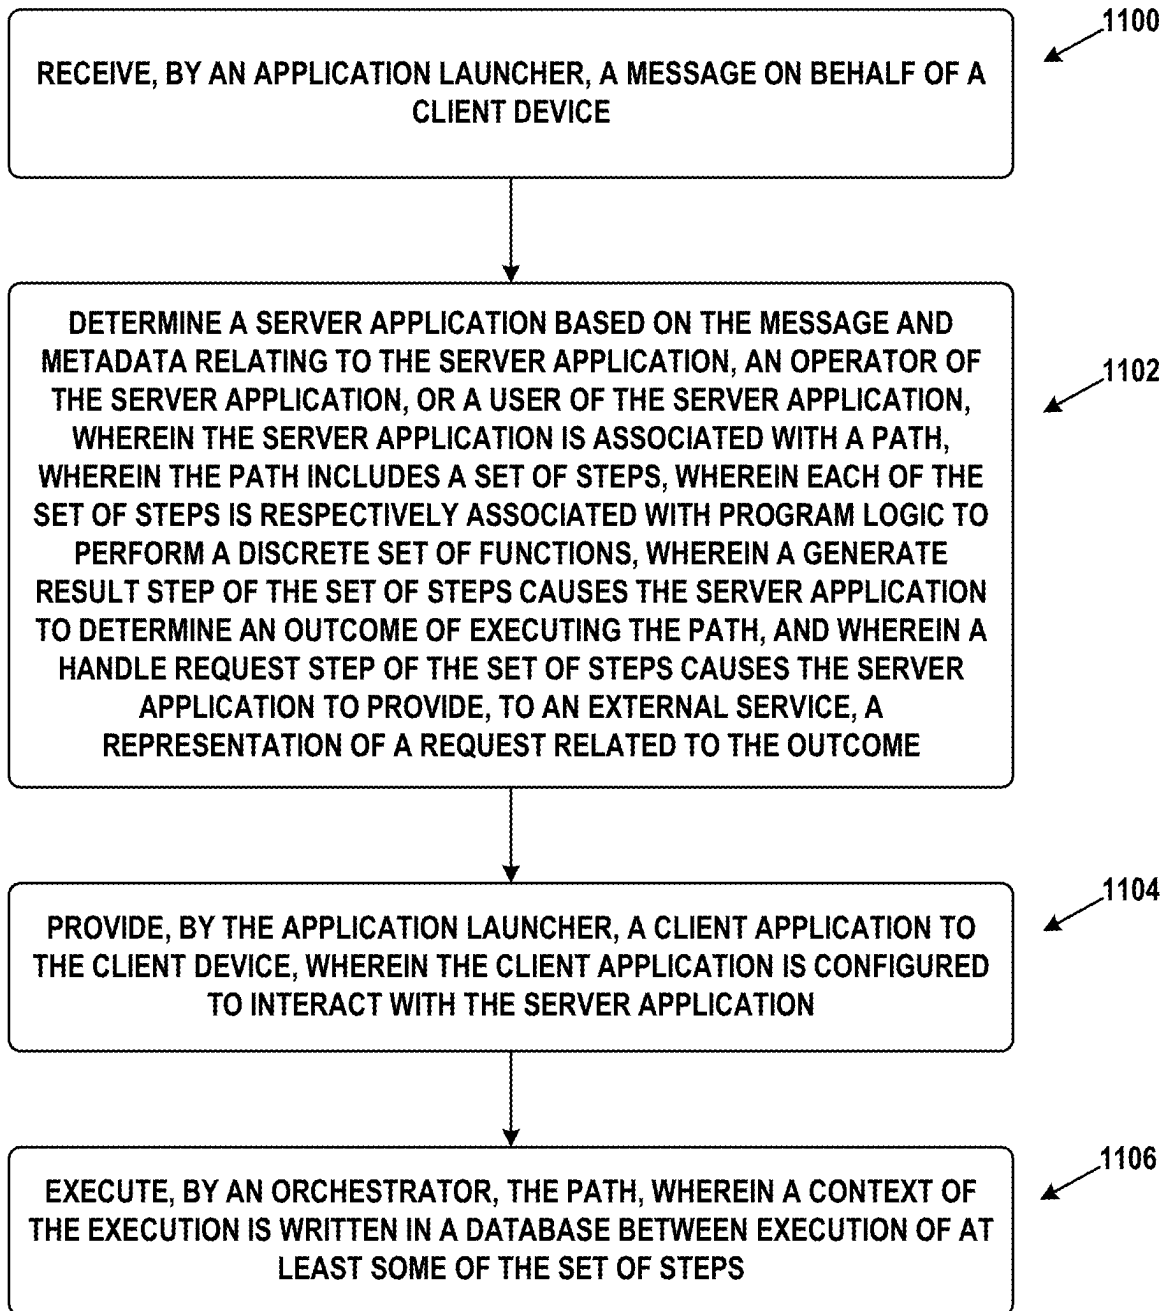


FIG. 11

TRANSACTION PROCESSING SYSTEM WITH DECOMPOSED ARCHITECTURE

CROSS-REFERENCE TO RELATED APPLICATION

[0001] This application is a continuation of and claims priority to U.S. patent application Ser. No. 18/523,161, filed Nov. 29, 2023, which is hereby incorporated by reference in its entirety.

BACKGROUND

[0002] A transaction path (or “path”) is a series of steps that are followed by a computing system in order to conduct a transaction. In an example embodiment, the transaction may include one or more discrete steps involving processing, communication with remote services, and/or storage of intermediate and final results. Existing systems that perform transactions are monolithic, rigidly architected, and overly complicated. As the technology has evolved, these existing systems were partially adapted for new types of transactions while retaining their monolithic nature. This evolution has accelerated in recent times, causing monolithic transaction paths to become brittle and error prone, with single points of failure. Making changes to these systems is time consuming, with extensive regression testing needed to ensure that the addition of new functionality does not inhibit the use or performance of pre-existing functionality.

SUMMARY

[0003] The embodiments herein provide technical improvements to these and potentially other technical problems by decomposing monolithic transaction processing systems into a number of discrete functional components. Each of these components is arranged to be able to perform a specific set of one or more functions. As a consequence, the decomposed server architecture is more flexible and reliable than its monolithic predecessors. Further, the decomposed server architecture can more readily be adapted to use parallelism, roll-back and roll-forward features, various pre-processing and post-processing steps, load balancing, and application loading time improvements. Other advantages are also possible.

[0004] Accordingly, a first example embodiment may involve one or more processors and memory storing instructions that are executable by the one or more processors to perform operations of software components comprising: a web site configured to allow selection of a server application, wherein the server application is associated with a path, wherein the path includes a set of steps, and wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions; storage containing metadata that relate to the server application and an operator of the web site; an application launcher configured to: receive a request on behalf of a client device, determine the server application based on the request and the metadata, and provide a client application to the client device, wherein the client application is configured to interact with the server application; and an orchestrator configured to execute the path, wherein at least some of the set of steps include interaction with one or more external services remotely accessible to the system, and wherein a context of the execution is written in a database between execution of at least some of the set of steps.

[0005] A second example embodiment may involve: receiving, by a web site, selection of a server application, wherein the server application is associated with a path, wherein the path includes a set of steps, and wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions; receiving, by an application launcher, a request on behalf of a client device; determining, by the application launcher, the server application based on the request, and also based on metadata that relate to the server application and an operator of the web site; providing, by the application launcher, a client application to the client device, wherein the client application is configured to interact with the server application; and executing, by an orchestrator, the path, wherein at least some of the set of steps include interaction with one or more external services, and wherein a context of the execution is written in a database between execution of at least some of the set of steps.

[0006] A third example embodiment may involve: one or more processors and memory storing instructions that are executable by the one or more processors to perform operations of software components comprising: a server application associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions, wherein a generate result step of the set of steps causes the server application to determine an outcome of executing the path, and wherein a handle request step of the set of steps causes the server application to provide, to an external service, a representation of a request related to the outcome; storage containing metadata that relate to the server application, an operator of the server application, or a user of the server application; an application launcher configured to: receive a message on behalf of a client device, determine the server application based on the message and the metadata, and provide a client application to the client device, wherein the client application is configured to interact with the server application; and an orchestrator configured to execute the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.

[0007] A fourth example embodiment may involve: receiving, by an application launcher, a message on behalf of a client device; determining a server application based on the message and metadata relating to the server application, an operator of the server application, or a user of the server application, wherein the server application is associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions, wherein a generate result step of the set of steps causes the server application to determine an outcome of executing the path, and wherein a handle request step of the set of steps causes the server application to provide, to an external service, a representation of a request related to the outcome; providing, by the application launcher, a client application to the client device, wherein the client application is configured to interact with the server application; and executing, by an orchestrator, the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.

[0008] A fifth example embodiment may involve a non-transitory computer-readable medium, having stored thereon program instructions that, upon execution by a computing

system, cause the computing system to perform operations in accordance with the first, second, third and/or fourth example embodiment.

[0009] In a sixth example embodiment, a system may include various means for carrying out each of the operations of the first, second, third and/or fourth example embodiment.

[0010] These, as well as other embodiments, aspects, advantages, and alternatives, will become apparent to those of ordinary skill in the art by reading the following detailed description, with reference where appropriate to the accompanying drawings. Further, this summary and other descriptions and figures provided herein are intended to illustrate embodiments by way of example only and, as such, that numerous variations are possible. For instance, structural elements and process steps can be rearranged, combined, distributed, eliminated, or otherwise changed, while remaining within the scope of the embodiments as claimed.

BRIEF DESCRIPTION OF THE DRAWINGS

[0011] FIG. 1 illustrates a schematic drawing of a computing device, in accordance with example embodiments.

[0012] FIG. 2 illustrates a schematic drawing of a server device cluster, in accordance with example embodiments.

[0013] FIG. 3 is a message flow diagram, in accordance with example embodiments.

[0014] FIG. 4 depicts a decomposed server architecture, in accordance with example embodiments.

[0015] FIG. 5 depicts modules representing steps that can be carried out by the decomposed server architecture, in accordance with example embodiments.

[0016] FIG. 6 depicts the saving of context between steps of a transaction path, in accordance with example embodiments.

[0017] FIG. 7 provides an example of a decomposed server architecture with load balancers, in accordance with example embodiments.

[0018] FIG. 8 depicts a decomposed web-based service architecture, in accordance with example embodiments.

[0019] FIG. 9 depicts a decomposed entertainment service architecture, in accordance with example embodiments.

[0020] FIG. 10 is a flow chart, in accordance with example embodiments.

[0021] FIG. 11 is a flow chart, in accordance with example embodiments.

DETAILED DESCRIPTION

[0022] Example methods, devices, and systems are described herein. It should be understood that the words “example” and “exemplary” are used herein to mean “serving as an example, instance, or illustration.” Any embodiment or feature described herein as being an “example” or “exemplary” is not necessarily to be construed as preferred or advantageous over other embodiments or features unless stated as such. Thus, other embodiments can be utilized and other changes can be made without departing from the scope of the subject matter presented herein.

[0023] Accordingly, the example embodiments described herein are not meant to be limiting. It will be readily understood that the aspects of the present disclosure, as generally described herein, and illustrated in the figures, can be arranged, substituted, combined, separated, and designed in a wide variety of different configurations. For example,

the separation of features into “client” and “server” components may occur in a number of ways.

[0024] Further, unless context suggests otherwise, the features illustrated in each of the figures may be used in combination with one another. Thus, the figures should be generally viewed as component aspects of one or more overall embodiments, with the understanding that not all illustrated features are necessary for each embodiment.

[0025] Additionally, any enumeration of elements, blocks, or steps in this specification or the claims is for purposes of clarity. Thus, such enumeration should not be interpreted to require or imply that these elements, blocks, or steps adhere to a particular arrangement or are carried out in a particular order.

I. Example Computing Devices and Cloud-Based Computing Environments

[0026] FIG. 1 is a simplified block diagram exemplifying a computing device 100, illustrating some of the components that could be included in a computing device arranged to operate in accordance with the embodiments herein. Computing device 100 could be a client device (e.g., a device actively operated by a user), a server device (e.g., a device that provides computational services to client devices), or some other type of computational platform. Some server devices may operate as client devices from time to time in order to perform particular operations, and some client devices may incorporate server features.

[0027] In this example, computing device 100 includes processor 102, memory 104, network interface 106, and input/output unit 108, all of which may be coupled by system bus 110 or a similar mechanism. In some embodiments, computing device 100 may include other components and/or peripheral devices (e.g., detachable storage, printers, and so on).

[0028] Processor 102 may be one or more of any type of computer processing element, such as a central processing unit (CPU), a co-processor (e.g., a mathematics, graphics, or encryption co-processor), a digital signal processor (DSP), a network processor, and/or a form of integrated circuit or controller that performs processor operations. In some cases, processor 102 may be one or more single-core processors. In other cases, processor 102 may be one or more multi-core processors with multiple independent processing units. Processor 102 may also include register memory for temporarily storing instructions being executed and related data, as well as cache memory for temporarily storing recently-used instructions and data.

[0029] Memory 104 may be any form of computer-usable memory, including but not limited to random access memory (RAM), read-only memory (ROM), and non-volatile memory (e.g., flash memory, hard disk drives, solid state drives, compact discs (CDs), digital video discs (DVDs), and/or tape storage). Thus, memory 104 represents both main memory units, as well as long-term storage. Other types of memory may include biological memory.

[0030] Memory 104 may store program instructions and/or data on which program instructions may operate. By way of example, memory 104 may store these program instructions on a non-transitory, computer-readable medium, such that the instructions are executable by processor 102 to carry out any of the methods, processes, or operations disclosed in this specification or the accompanying drawings.

[0031] As shown in FIG. 1, memory 104 may include firmware 104A, kernel 104B, and/or applications 104C. Firmware 104A may be program code used to boot or otherwise initiate some or all of computing device 100. Kernel 104B may be an operating system, including modules for memory management, scheduling and management of processes, input/output, and communication. Kernel 104B may also include device drivers that allow the operating system to communicate with the hardware modules (e.g., memory units, networking interfaces, ports, and buses) of computing device 100. Applications 104C may be one or more user-space software programs, such as web browsers or email clients, as well as any software libraries used by these programs. Memory 104 may also store data used by these and other programs and applications.

[0032] Network interface 106 may take the form of one or more wireline interfaces, such as Ethernet (e.g., Fast Ethernet, Gigabit Ethernet, and so on). Network interface 106 may also support communication over one or more non-Ethernet local-area media, such as coaxial cables or power lines, or over wide-area media, such as fiber-optic connections (e.g., OC-x interfaces) or digital subscriber line (DSL) technologies. Network interface 106 may additionally take the form of one or more wireless interfaces, such as IEEE 802.11 (Wifi), Bluetooth, global positioning system (GPS), or a wide-area wireless interface (e.g., using 4G or 5G cellular networks). However, other forms of physical layer interfaces and other types of standard or proprietary communication protocols may be used over network interface 106. Furthermore, network interface 106 may comprise multiple physical interfaces. For instance, some embodiments of computing device 100 may include Ethernet, Bluetooth, and Wifi interfaces.

[0033] Input/output unit 108 may facilitate user and peripheral device interaction with computing device 100. Input/output unit 108 may include one or more types of input devices, such as a keyboard, a mouse, a touch screen, and so on. Similarly, input/output unit 108 may include one or more types of output devices, such as a screen, monitor, printer, and/or one or more light emitting diodes (LEDs). Additionally or alternatively, computing device 100 may communicate with other devices using a universal serial bus (USB) or high-definition multimedia interface (HDMI) port interface, for example.

[0034] In some embodiments, one or more computing devices like computing device 100 may be deployed as a cluster of server devices. The exact physical location, connectivity, and configuration of these computing devices may be unknown and/or unimportant to client devices. Accordingly, the computing devices may be referred to as “cloud-based” devices that may be housed at various remote data center locations.

[0035] FIG. 2 depicts a cloud-based server cluster 200 in accordance with example embodiments. In FIG. 2, operations of a computing device (e.g., computing device 100) may be distributed between server devices 202, data storage 204, and routers 206, all of which may be connected by local cluster network 208. The number of server devices 202, data storages 204, and routers 206 in server cluster 200 may depend on the computing task(s) and/or applications assigned to server cluster 200.

[0036] For example, server devices 202 can be configured to perform various computing tasks of computing device 100. Thus, computing tasks can be distributed among one or

more of server devices 202. To the extent that these computing tasks can be performed in parallel, such a distribution of tasks may reduce the total time to complete these tasks and return a result. For purposes of simplicity, both server cluster 200 and individual server devices 202 may be referred to as a “server device.” This nomenclature should be understood to imply that one or more distinct server devices, data storage devices, and cluster routers may be involved in server device operations.

[0037] Data storage 204 may be data storage arrays that include drive array controllers configured to manage read and write access to groups of hard disk drives and/or solid state drives. The drive array controllers, alone or in conjunction with server devices 202, may also be configured to manage backup or redundant copies of the data stored in data storage 204 to protect against drive failures or other types of failures that prevent one or more of server devices 202 from accessing units of data storage 204. Other types of memory aside from drives may be used.

[0038] Routers 206 may include networking equipment configured to provide internal and external communications for server cluster 200. For example, routers 206 may include one or more packet-switching and/or routing devices (including switches and/or gateways) configured to provide (i) network communications between server devices 202 and data storage 204 via local cluster network 208, and/or (ii) network communications between server cluster 200 and other devices via communication link 210 to network 212.

[0039] Additionally, the configuration of routers 206 can be based at least in part on the data communication requirements of server devices 202 and data storage 204, the latency and throughput of the local cluster network 208, the latency, throughput, and cost of communication link 210, and/or other factors that may contribute to the cost, speed, fault-tolerance, resiliency, efficiency, and/or other design goals of the system architecture.

[0040] As a possible example, data storage 204 may include any form of database, such as a structured query language (SQL) database. Various types of data structures may store the information in such a database, including but not limited to tables, arrays, lists, trees, and tuples. Furthermore, any databases in data storage 204 may be monolithic or distributed across multiple physical devices.

[0041] Server devices 202 may be configured to transmit data to and receive data from data storage 204. This transmission and retrieval may take the form of SQL queries or other types of database queries, and the output of such queries, respectively. Additional text, images, video, and/or audio may be included as well. Furthermore, server devices 202 may organize the received data into web page or web application representations. Such a representation may take the form of a markup language, such as the HyperText Markup Language (HTML), the extensible Markup Language (XML), Cascading Style Sheets (CSS), and/or JavaScript Object Notation (JSON), or some other standardized or proprietary format. Moreover, server devices 202 may have the capability of executing various types of computerized scripting languages, such as but not limited to Perl, Python, PHP Hypertext Preprocessor (PHP), Active Server Pages (ASP), JavaScript, and so on. Computer program code written in these languages may facilitate the providing of web pages to client devices, as well as client device interaction with the web pages. Alternatively or

additionally, Java may be used to facilitate generation of web pages and/or to provide web application functionality.

II. Example Transaction System

[0042] FIG. 3 depicts an example monolithic transaction processing system that can be used to facilitate various types of interactions with client devices (e.g., enabling web, email, storage, streaming, and/or other services). Such a monolithic system has a number of drawbacks relating to complexity, defect rates, testing, adaptability, and so on. Nonetheless, it can be instructive to consider how such a monolithic system operates so that the improvements provided by the embodiments herein can be fully appreciated.

[0043] FIG. 3 includes client device 300, server device 302, database 304, external service 306, and external service 308. These components may be communicatively coupled with one another such that client device 300 can communicate with server device 302, and server device 302 can communicate with each of database 304, external service 306, and external service 308. But other arrangements are possible.

[0044] Client device 300 may be a computing device as described in the context of FIG. 1. Thus, client device 300 may be a desktop computer, laptop computer, mobile phone, tablet, etc. To interact with server device 302, client device 300 may execute one or more dedicated client applications. Alternatively, client device 300 may do so by way of a web browser application or another form of web-based interface. In some environments, there may be many instances of various types of client devices (e.g., tens, hundreds, or thousands of client devices).

[0045] Server device 302 may be a computing device or computing cluster as described in the context of FIGS. 1 and/or 2. Thus, server device 302 may include one or more physical server devices in various arrangements. To interact with other devices, server device 302 may execute one or more server applications. These server applications may be multithreaded, capable of executing multiple instances simultaneously, and/or distributed across different units of physical hardware.

[0046] For sake of convenience, even where multiple server applications exist, they will be collectively referred to as “the server application” herein. In general, references to server device 302 and the server application may be used interchangeably unless context indicates otherwise.

[0047] Database 304 may provide an organized collection of data that is accessible to server device 302. Database 304 may operate on a physical computing device that is distinct from that of server device 302 or it may be incorporated into server device 302 in some fashion. Thus, database 304 may be purely software or a combination of hardware and software. Regardless, database 304 may support a number of operations, such as storing, providing, indexing, searching, querying, and updating data. In various embodiments, database 304 may operate as a relational database (e.g., using SQL), a NoSQL database (e.g., MongoDB), an in-memory database (e.g., Redis), or a graph database (e.g., Neo4j). In a relational database, for example, this data may be stored in one or more tables in accordance with a predefined schema. Other possibilities exist. In some embodiments, database 304 may represent multiple physical or logical databases. Each of these may be replications of one another for purposes of redundancy and/or load balancing, or each may have a different intended use.

[0048] External services 306 and 308 may be remote services accessed by way of respectively-defined application programming interfaces (APIs). These APIs may be based on representational state transfer (REST) interfaces, or take some other form such as remote procedure call (RPC), simple object access protocol (SOAP), or Websockets. Thus, external services 306 and 308 may be authentication services (used to authenticate users of client device 300 that access server device 302), geolocation services (used to identify the physical location of client device 300 and/or a user of client device 300), machine learning services (used to apply trained machine learning models to classify the provided input), generative artificial intelligence (AI) services (to generate text, images, audio, and/or video based on a prompt), and/or additional database services (used to provide and/or store data related to client device 300 and/or server device 302). Alternatively or additionally, external services 306 and 308 may provide various types of application-specific services, such as web, email, social media, or streaming, services.

[0049] In some environments, one or more of external services 306 and 308 may be provided by third parties that are distinct entities from those that own or operate client device 300, server device 302, or database 304. In these environments, the third parties may be referred to as “operators.” In other environments, one or more of external services 306 and 308 may be integrated with server device 302 and/or database 304. Thus, use of the term “external” does not require that external services 306 and 308 are provided by a third party, and instead may be provided by the entity that operates server device 302. For example, external services 306 and 308 may be “inline” services performed by server device 302 and/or database 304.

[0050] In general, FIG. 3 depicts messages flowing between elements of a transaction processing system. Such a system may refer to just server device 302 and its operations, or server device 302, database 304, any of the external services, and their combined operations. In this figure, communicated information generally flows in the direction of the arrows. However, even steps represented by unidirectional arrows may involve at least some information flowing in the opposite direction (e.g., acknowledgements to messages transmitted in the direction of an arrow).

[0051] As shown in FIG. 3, a transaction may include a number of steps known as a transaction path (or simply, a “path”), and these steps may be carried out by server device 302 in a particular order. Different transaction paths may carry out different combinations of steps in different orderings, and all of these transaction paths may be monolithically implemented in the server application. Here, such a monolithic implementation means that the transaction paths are explicitly coded within a single large application or a small number of large applications (e.g., all components of these applications are interwoven and tightly coupled, often into a single program on a single platform).

[0052] As one way of initiating a transaction, client device 300 may engage in step 310 with server device 302. Step 310 may be transmission of a request, query, command, API call, or other form of invocation that can cause server device 302 to at least attempt to carry out a specific step of the transaction path. The type of step may be implied or made explicit in step 310. In response to step 310, server device 302 may carry out steps 312, 314, 316, 318, and 320. In various embodiments, more or fewer steps may be carried out.

[0053] Step 312 may involve server device 302 communicatively interacting with database 304 to obtain application state relating to execution of the application. In some cases, this application state is specific to a client device, such as client device 300. The communication may include server device 302 querying database 304 to read and obtain this state without making any updates to database 304. Thus, step 312 may be a read-only operation. The state may contain sufficient information (e.g., application configuration data) so that server device 302 can perform the remaining steps of the transaction path. Alternatively, server device 302 may obtain further application state from database 304 at other points in the transaction path (not shown).

[0054] Step 314 may involve server device 302 communicatively interacting with external service 306 to generate provisional results to add to the application state. Doing so may include server device 302 providing at least some of the application state obtained in step 312 to external service 306 by way of a communication interface. Doing so may also include server device 302 receiving the provisional results from external service 306, again by way of a defined interface. In some cases, server device 302 may pre-process the application state before providing it to external service 306, post-process the provisional results before adding them to the application state, and/or generate part of the application state.

[0055] In some environments, external service 306 may be an authentication service that determines whether client device 300 is permitted to log into or obtain further information from server device 302. In these environments, the application state may include credentials for a user of client device 300, such as a userid and password. In other environments, external service 306 may provide results of execution of an application such as a machine learning classifier, a generative AI program, or some other type of application.

[0056] Step 316 may involve server device 302 writing the results from step 314 and any subsequent post-processing to database 304 as part of the application state. As shown in FIG. 3, what is written may be referred to as the provisional state, in that it represents the outcome of an intermediate step. Step 316 may take place immediately or shortly after step 314 concludes and may occur in response to step 314.

[0057] Step 318 may involve server device 302 communicatively interacting with external service 308 to generate further results to add to the application state. Doing so may include server device 302 providing at least some of the application state obtained in steps 312 and/or 314 to external service 308 by way of a defined interface. Doing so may also include server device 302 receiving the further results from external service 308, again by way of a defined interface. In some cases, server device 302 may pre-process the application state as it stands after step 314 before providing it to external service 308, post-process the further results before adding them to the application state, and/or generate part of the application state. This pre-processing and post-processing may involve transforming the results and/or application state into a consistent format or making other modifications for purposes of facilitating storage and processing.

[0058] In some environments, external service 308 may be one of various types of application services. The invocation of this application service may be a consequence of application state obtained in steps 312 and/or 314. For example, external service 308 may be a web-based service (arranged to provide online web-based information to server device

302), an email service (arranged to provide interaction with email accounts to a user of client device 300), a social media platform (arranged to provide social media posts, updates, messages, comments, and/or news feeds to a user of client device 300), a storage system (arranged to store data and provide copies of stored data on behalf of client device 300 or server device 302), or a streaming service (arranged to transmit a stream of audio and/or video to server device 302 and/or client device 300).

[0059] Step 320 may involve server device 302 writing the further results from step 318 and any subsequent post-processing to database 304 as part of the application state. As shown in FIG. 3, what is written may be referred to as the further state, in that it represents the outcome of further step of the transaction path. In some cases, the further state may be a final state in that it represents the final outcome of the carrying out the transaction path. Step 320 may take place immediately or shortly after step 318 concludes and may occur in response to step 318.

[0060] Step 322 may involve server device 302 transmitting a response to client device 300 containing a representation of the outcome of the transaction. This outcome may be as simple as a Boolean indicator of success or failure (e.g., true or false, respectively) or instructions or program code that cause a user interface on client device 300 to display the outcome in some manner. This may involve client device 300 displaying some form of structured information (e.g., dashboards, tables, or charts), images, and/or animations.

[0061] The transaction processing system depicted in FIG. 3 can be implemented in a stateless fashion. Thus, the steps performed by server device 302 are not dependent on any in-memory (e.g., in RAM) state during their execution. This stateless design means that if the hardware of server device 302 was to turn off and/or all in-memory state was lost at any point in time, server device 302 could restart and gracefully resume from where it left off. The actual state of the transaction is always stored in database 304.

[0062] Notably, step 312 does not change the state as stored in database 304. Step 316 saves the provisional state to database 304 with the understanding that this provisional state has not yet been used in step 318. Step 320 saves the further (perhaps final) state to database 304 with the understanding that this further state has not yet been used in step 322. In the case that step 318 can either succeed or fail, the final state may be written to database 304 when step 318 is successful, but the in-memory representation of this state may be marked as void or invalid when step 318 fails.

[0063] Given this software architecture of server device 302, execution of the transaction path may occur multiple times in parallel client sessions for different client devices or different client applications on client devices. It will only use local variables dedicated to those client sessions, so there is no need for explicit mutual exclusion mechanisms (e.g., mutexes, spinlocks, or semaphores) for avoidance of race conditions.

[0064] This design allows many requests to be processed in parallel. Since the transactions are stateless, a number of server application instances can execute simultaneously that is only limited by the processing, storage, and communications limits of the transaction processing system. Typically, database 304 is the convergence point where multiple transactions can be in various stages the same time. It is possible

to split users and/or client devices across multiple instances of database 304 to distribute transactions and improve performance.

[0065] Nonetheless, the transaction processing system design of FIG. 3 does have some drawbacks. Particularly, it does not easily or efficiently support new features or functionalities. A consequence is that development, maintenance, and support of this software design can grow complex and cumbersome.

[0066] As one example, step 312 can be thought of as what organizes the relevant information that can be used in step 314 that follows. Primarily, step 312 involves retrieving, from database 304, the most up to date version of the application state so that this application state can be used to generate the provisional result. The step of organizing the relevant information has grown in scope based on evolving security requirements, application capabilities, and user differentiations. Thus, there have been more and more checks put into place regarding how a user is permitted to interact with the server application.

[0067] For instance, user capabilities may be geo-fenced so that users in different geographic locations (e.g., states, provinces, regions, or countries) have different degrees of access to server application features. Furthermore, the server application may have controls over how many transactions, the types of transaction, and/or the magnitude of each transaction in which a user may participate (here, transaction magnitude may involve an amount of data exchanged between client device 300 and server device 302, the amount of processing server device 302 requires to carry out the transaction, the amount of time that the user engages with the server application, and/or a pecuniary value of the transaction). Moreover, certain users might have priority over others in some fashion so that their transactions are processed more expeditiously. Given these requirements, the server application may also check the integrity of some or all incoming data packets in the request of step 310 to ensure that the data packets have not been illicitly modified (i.e., hacked).

[0068] In any event, transactions that otherwise would have been triggered by step 312 may be blocked by server device 302 if they do not meet any of these requirements. This means that server device 302 and/or database 304 should understand and support the relevant requirements, check these requirements, and then act accordingly. This increases the complexity of server device 302 and/or database 304. Notably, any change to one of these features or the addition of a new feature requires updates to the server application as well as database 304.

[0069] As another example, step 314 is designed to treat external service 306 as a stateless “black box.” As such, external service 306 receives all of its input in a request or query transmitted by server device 302 and provides intermediate results based on this input as well as its programming. Notably, the server application and external service 306 may be developed independently by different entities, so communication therebetween by way of a strictly-defined interface may be required. The server application cannot change the input and expect external service 306 to be configured to understand this new input arrangement. Likewise, external service 306 should provide to the server application the intermediate result in a form that the server

application is configured to understand. The same constraints apply to interactions between server device 302 and external service 308.

[0070] By nature of database 304 being the sole source of truth for all transaction data, external service 306 might not have access to any other source of input and might not save any state associated with the server application. What this amounts to is that when a new feature or functionality is developed, both the server application and external service may need to be modified to support that feature or functionality. Again, the same or similar limitations may apply to interactions between server device 302 and external service 308.

III. Improved Transaction Processing System

[0071] To modify the transaction processing system of FIG. 3 so that it is more flexible and extendible, the programming logic of server device 302 and database 304 can be changed. In some situations, it is desirable to only modify server device 302 and database 304 to support new steps and transactions. While it is possible to modify the external services, doing so may require that these services are also modified to support synchronization, rollbacks, reporting, coordination, and orchestration in order to align their operations with those of server device 302 and database 304. Therefore, it is simpler to limit the modifications to server device 302 and database 304 where possible. For example, modifications may be made to server device 302 and database 304 as well as the interfaces between server device 302 and the external services in order to reduce the impact on the external services.

[0072] Furthermore, some of the new steps and transactions may involve integrating new technologies into transaction paths. For example, these can be based on so-called web3 technology that supports blockchain to distribute immutable database records amongst multiple parties. Using blockchain, these parties can obtain consensus regarding transactions or other behaviors of users or entities. Blockchain mechanisms can include smart contracts and cryptographic tokens that represent various capabilities with respect to the server application and/or external services. For sake of example, U.S. Patent Application Publication No. 2023/0315819 and PCT Patent Application Publication No. WO2023193971, both hereby incorporated by reference herein, describe how non-fungible tokens (NFTs) can be used to control access to various features of software applications. Integrating these new technologies into current systems is simply not practical, as a re-engineering of the server application and database 304 is required for each.

[0073] Moreover, at least some of the external services may take the form of microservices. Microservices can be part of decomposed architectures in which an application includes or relies upon a set of relatively small and discrete networked or integrated services. Thus, the microservices architecture stand in contrast to a monolithic application in that each of the microservices may be independently maintainable, testable, and deployable and/or otherwise loosely coupled to one another and the overall application.

[0074] For example, a web site that allows authenticated users to have access to a particular application may be designed as a server application that makes use of microservices for: (i) processing HyperText Transfer Protocol (HTTP) requests and providing appropriate static and/or dynamic web-based content (e.g., HTML, CSS, JavaScript)

in corresponding HTTP responses, (ii) allowing new users to register for the service and authenticating existing users as they log on, (iii) session management for logged in users, and (iv) one or more database interfaces to access stored web content. Other possibilities exist. Notably, many different web sites can be based on this or a similar arrangement of microservices.

[0075] As another example, an external service may provide streaming of video content to registered users may be designed as a server application that makes use of microservices for: (i) processing user requests for video content and providing the requested video content in response to these requests, (ii) allowing new users to register for the service and authenticating existing users as they log on, (iii) session management for logged in users, (iv) one or more database interfaces to access stored video content, (v) a search, browse, and/or recommendation service for users to find video content based on metadata like titles, descriptions, genres, and actors, (vi) a monitoring and analytics service that collects data on user engagement, streaming quality, and service performance, and (vii) a content delivery network gateway that provides mechanisms for distributing video content across geographical locations to reduce latency experienced by users. Other possibilities exist. Notably, many different streaming services can be based on this or a similar arrangement of microservices.

[0076] Additionally, some microservices may be common and/or generalizable between these video streaming services and the aforementioned web site services (e.g., registration, authentication, session management, etc.). Thus, common units of software that perform and/or interact with these microservices may be used across different transaction paths within multiple applications of the same or a different type.

[0077] Calls to these microservices (e.g., the server application invoking one or more of the microservices) may be arranged in various ways and orderings. In some cases, certain sets of calls may be performed in parallel. In other cases, a particular sequence of calls might be performed serially. For a given transaction path, some combination of parallel and serialized calls may be possible.

[0078] The employment of microservices in the manner described herein allows a server architecture to be able to perform one more or more distinct sets of server application functions (e.g., web-based services, streaming services, email services, machine-learning and AI services) in a modular and efficient manner. Doing so avoids two disadvantageous server application designs of previous transaction processing systems. The first is to place all desired functionality into a single server application, resulting in a complex, unwieldy, and fragile application that is difficult to test, debug, maintain, and expand. The second is to have different server applications for each transaction path, resulting in the need to design, maintain, and test the overlapping functionality of each server application independently.

[0079] Regarding terminology, a transaction path is an arrangement of steps performed by one or more server devices of the transaction processing system, perhaps with assistance from external services. A server application may support the execution of one or more transaction paths. A transaction is a discrete execution of the steps of a transaction path, such that the transaction has a beginning and an end as defined by the transaction path. Thus, a server application might support multiple transaction paths, each of

which may be used to carry out many transactions. Nonetheless, these terms are intended to be used in a flexible and non-limiting fashion.

A. Decomposed Server Architecture

[0080] FIG. 4 depicts a decomposed server architecture **400** that provides numerous advantages over previous transaction processing systems. These advantages include but are not limited to reuses of modular software components, less need for testing, and a more robust and extensible design. Consequently, use of server architecture **400** results in higher performance, greater stability, and more flexibility than the previous approaches.

[0081] In FIG. 4, objects with solid lines generally refer to software modules and objects with dashed lines generally refer to data. However, some software modules (e.g., steps **414**) may be treated as data.

[0082] Server architecture **400** includes software modules for operator web site **402**, application launcher **404**, location/geoblocking service **406**, application metadata **408**, operator metadata **410**, and orchestrator **416**. Based on this architecture, a user of client application **300A** on client device **300** can request use of a particular “web site” or “application”, and application launcher **404** can determine the appropriate one or more of transaction paths **412**. Each of transaction paths **412** is defined at least in part by one or more of steps **414**. Orchestrator **416** may contain or have access to database **418** that may maintain initial, intermediate, and/or final state for steps **414** as orchestrator **416** carries out these steps. Further, some of steps **414** may interact with operator bridge **422** to access operator services **424** and/or interact with transaction database **426** to provide information for monitoring services **428**.

[0083] Notably, server architecture **400** may contain more or fewer software modules arranged in various ways. Thus, server architecture **400** is just one example of how monolithic server applications can be decomposed and made more flexible. Each of these software modules is briefly described below.

[0084] For sake of convenience, even where multiple server applications exist, they will be collectively referred to as “the server application” herein. In general, references to server architecture **400** and the server application may be used interchangeably unless context indicates otherwise.

[0085] Operator web site **402** may be a web-based interface and/or portal through which a user can register, login, edit their account settings and preferences, and so on. There may be many different instances of operator web site **402** (e.g., for different operators and accessible by different URLs) supported by server architecture **400**. Here, operators may refer to entities that provide, control, and/or manage the applications available to client devices. Some operators may require that similar application behave in different ways from other operators.

[0086] Application launcher **404** may process the launching of a specific application defined by one or more transaction paths. The application may be implicitly associated with the selected operator web site **402** or the user may select the specific application by way of operator web site **402**. Application launcher **404** may provide client application **300A** (e.g., by way of operator web site **402**) to client device **300**. Client application **300A** may be specific to the application that was selected and may executed in a web browser of client device **300** (e.g., as an HTML5 applica-

tion) or may be a standalone application (or “app”). Thus, application launcher **404** may store or have access to one or more instances of executable code or executable software image files per application for execution on a client device (here, an executable software image file contains program code capable of being executed on a computer).

[0087] In order to launch an application, application launcher **404** may access one or more of location/geoblocking service **406**, application metadata **408**, and/or operator metadata **410** to determine the version of the application to launch, its operational environment and parameters, and which it can be launched by client device **300**. A successful launch will result in one or more transaction paths **412** being provided to orchestrator **416**. A given transaction path may include one or more of steps **414**, where each of these steps is a discrete set of functionality.

[0088] Location/geoblocking service **406** may be a service that maps client devices, their IP addresses, users, and/or the reported locations of any of these entities to a geographic region. In cases where an application or transaction path is limited to certain regions or is configured to be blocked in other regions, location/geoblocking service **406** may transmit sufficient information to an edge device, router or firewall so that the client device in question can or cannot use the application. For example, certain applications may only be available in particular locations due to language, cultural, regulatory, or other concerns.

[0089] Application metadata **408** may store metadata about each application (e.g., the application’s name, operator, allowed or blocked regions, etc.). On an attempted application launch, application metadata **408** can be checked to determine whether the client device is permitted to access the application.

[0090] Operator metadata **410** may store metadata about an operator, such as an operator name, setting, and general environment (e.g., graphical themes). On an attempted application launch, operator metadata **410** can be checked and/or accessed to control the operational parameters of the application.

[0091] As noted above, orchestrator **416** may receive representation of one or more transaction paths **412**. For sake of simplicity in the discussion below, it is assumed herein that just one transaction path **412** is provided. This transaction path **412** may include one or more steps **414**. The transaction path **414** may also define one or more orderings of steps **414** as well as indications of which of these steps may be performed in parallel and which must be performed serially. For example, the transaction path **414** may indicate which steps should follow other steps of steps **414**. In general, steps **414** may be carried out one or more times by orchestrator **416** per transaction. Thus, a transaction path **412** may also include some indication of a loop or repetition amongst steps.

[0092] Given all of this, a transaction path **412** may represent steps **414** in the form of a directed graph. Each of steps **414** may be a node in such a graph and an edge directed (e.g., with an arrow) from a first node to a second node may indicate that the step represented by second node is to be executed after the step represented by the first node. But other possibilities exist.

[0093] External services **420** may be a set of services that are hosted remotely or outside of server architecture **400** or

the transaction processing system. External services **420** may be invoked or accessed (e.g., by a defined API) by one or more of steps **414**.

[0094] For example, one or more of steps **414** may access operator bridge **422**, which services as a gateway, proxy, or interface to operator services **424**. Operator services **424** may include functionality that manages the relationship of client device **300** and/or its user with a particular operator. Where there are multiple operators, there may be multiple instances of operator bridge **422** and operator services **424** (e.g., one of each for every operator).

[0095] In another example, one or more of steps **414** may access transaction database **426**. For instance, one or more of steps **414** may transmit log data regarding their execution and/or its results to transaction database **426**. Transaction database **426** could be an SQL or NoSQL database, or take some other form. Monitoring services **428** may read this log data to determine performance, correctness, and other aspects of the server application. In some cases, monitoring services **428** may be managed and/or operated by a third party that is responsible for overseeing the performance of the server application.

[0096] Thus, in some cases, each of steps **414** may access an external service. Alternatively, at least some of steps **414** may be in-line (i.e., executed by the server application without invoking an external service).

B. Step Properties

[0097] FIG. 5 provides more examples of steps **414**. Particularly, steps **414** may be any one of steps **500-534** shown in FIG. 5, though other types of steps are possible. As noted, different arrangements and/or subsets of steps **414** may be represented by transaction paths **412**.

[0098] Session management **500** may involve retrieving and maintaining session information for the transaction. This could include parts of application metadata **408**, operator metadata **410**, and/or user-specific information.

[0099] Retrieve application state **502** may involve obtaining (e.g., from database **418** or elsewhere) the application state and/or the user’s progress within the server application as represented by steps of the transaction path.

[0100] Establish transaction ID **504** may involve generating a transaction identifier (e.g., a unique number or token) representing the current transaction.

[0101] Get application settings **506** may involve obtaining operational settings for the server application (e.g., from application metadata **408**, operator metadata **410**, and/or user-specific information).

[0102] Handle request **508** may involve receiving and processing a request from client application **300A**, as well as possibly providing a representation of the request to the operator by way of operator bridge **422** and operator service **424**.

[0103] Generate result **510** may involve determining a result (e.g., in response to a request received in handle request **508**) and possibly providing the result to client application **300A**. This result may be determined locally by the server application or in conjunction with the operator by way of operator bridge **422** and operator service **424**.

[0104] Finalize transaction **512** may involve storing the result of the transaction to database **418**, as well as possibly providing a representation of the result to the operator by way of operator bridge **422** and operator service **424**.

[0105] End transaction **514** may involve completing and ending the transaction. For example, this may involve deallocating memory used for the transaction and otherwise marking the transaction as complete, as well as possibly notifying the operator that the transaction is complete by way of operator bridge **422** and operator service **424**.

[0106] Log transaction **516** may involve publishing data related to the transaction to transaction database **426**. This data may include, for example, timestamps of times at which various steps of the transaction took place, the names or other representations of the steps, the user, server application, and/or operator involved, statistics about the transaction (e.g., length in time, amounts transacted, options selected by the user, features of the server application used, and so on). The published data may be stored in transaction database **426** possibly in a manner accessible to monitoring services **428**.

[0107] In addition, each of steps **500-516** may be subject to pre-step processing **520** and/or post-step processing **522**. In other words, some of steps **500-516** may be preceded by an instance of pre-step processing **520**, some of steps **500-516** may be followed by an instance of post-step processing **522**, and some of steps **500-516** may be preceded by an instance of pre-step processing **520** and an instance of post-step processing **522**. This pre-step and post-step processing may be different for some of steps **500-516**. Alternatively, such processing may be common for some of steps **500-516**. As an example, one or more of steps **500-516** may be associated with a pre-processing step that saves application state prior to the step being carried out and/or a post-processing step that saves application state after the step is carried out.

[0108] Put another way, orchestrator **416** may function as a scheduler, facilitating the performance of steps **414** required for a specific transaction path **412**. It initiates a sequence of steps **414** at the required times. To do so, orchestrator **416** may receive a request from client application **300A** and responsively execute a defined transaction path **412** through steps **414** based on the request. Depending on the server application, there might be need for different combinations of steps **414** to be in the transaction path **412**.

[0109] Get test data **530**, perform test **532**, and save test state **534** allow each of steps **414** and/or sequences thereof to be tested independently of the rest of the steps **414** in transaction path **412**. Advantageously, this allows test data to be used to instrument the application state, perform one or more tests on this instrumented application state, and save the results of these tests. These built-in testing steps permit the construction of application states that are expected to be rare or at least improbable in real-world deployments.

[0110] The sequence of steps **414** allows for some of these steps to be performed in parallel, while other steps are performed sequentially. The transaction path **412** can be customized to the nature of the request, containing as many of steps **414** as required in various orderings. Each of steps **414** may execute some program logic, and at least part of this execution can be carried out via a call to an external service.

C. Saving of Context and Step Independence

[0111] When execution of a transaction path **412** is initiated, an entry is created in a database **418** for the execution. At any point during the processing of the transaction path **412**, the state of the server application with respect to the

transaction path **412** can be saved into this entry in database **418**. As noted, the transaction path **412** allows for configuration of each of steps **414**, including the ability to save the current application state before the step is executed and after the step has executed. The application state stored in database **418** is also referred to as “context.” This context can be passed into each of steps **414** in the transaction path **412** and the respective step can use this information, as well as update, edit, and/or add to this information. Thus, the context may change as execution of the transaction path **412** progresses.

[0112] An example of data can be added to the context and used in a later step is as follows. Any of steps retrieve application state **502**, establish transaction ID **504**, and get application settings **506** may update the context. When the generate result **510** step is called, it may use what was stored in the context in order to determine the result. Thus, various instances of steps **414** may add to the context and use it as required. A tight coupling of the data used between these steps is not required. The data being used in the transaction processing is highly flexible, the context data structure is not bound to any schema of database **418** or steps being executed.

[0113] By saving the context while executing the transaction path **412**, it facilitates a degree of fault tolerance. If execution of the transaction path is halted due to a powering off of a device, an application crash, a lack of network connectivity, or some other hardware or software fault, the exact state of the transaction path **412** has been saved. When the transaction path **412** restarts (e.g., after a reboot of hardware, relaunch of one or more applications, or a regaining of network connectivity), orchestrator **416** can read this state and determine how to proceed. For example, orchestrator **416** could continue to process the transaction path **412** from where it left off, or determine that it needs to undo something that has been done and update the context accordingly.

[0114] FIG. 6 depicts the saving of context between steps. Notably, flow diagram **600** includes steps **602**, **606**, and **610**, which may be selected from steps **414** and which may be part of a transaction path **412**. Save context operation **604** optionally occurs between steps **602** and **606**. Similarly, save context operation **608** optionally occurs between steps **606** and **610**.

[0115] Each of steps **414** may be responsible for its execution independently from any other of steps **414**. In this manner, steps **414** do not have interaction with the transaction path **412** or other of steps **414**—instead, a step is passed the current context and executes with this context. It may happen that the connection to a particular step’s service is lost and this service needs to be called again in order for the step to complete. Such a step is designed so that it may be called twice or more with the same context, but it will only perform the execution once.

[0116] In the case that a step is stateless (e.g., the step itself does not maintain any state of its own), then the execution can easily be called again. If the step stores any data or calls to any external service, it is designed to be responsible for being called again and for handling second and subsequent calls correctly. In other words, the step may exhibit idempotency, in that it guarantees that performing the step multiple times will yield the same result as if it was executed only once.

[0117] Once a step has executed successfully, it can add data to the context. This data can then be saved after the current step and before the next step is executed. This will be saved to database 418. Thus, the step can provide, by way of the context, data in a completed state to the next step in the transaction path. This next step may also use the context.

[0118] It should be noted that a step might not be able to communicate directly with another step. Instead, steps communicate with one another by storing data in and reading data from the context. This principle helps maintain the transaction processing software architecture as a composition of small, well-defined, and independent pieces of logic. Therefore, future changes can be made without impacting large pieces of code and other logic. Indeed, individual steps can be modified, upgraded, adapted or changed often without affecting the other steps in the transaction path. This independence contrasts with existing systems, where all logic is collectively hosted and changes are exceedingly complicated. The independence of steps also allows people to modify/work on steps at the same time, without needing to modify a single file. This adds flexibility when adding new applications types, as new types of microservices can be created, tested, and modified outside of the existing transaction path.

D. Step Parallelization and Serialization

[0119] As noted above, transaction path 412 and orchestrator 416 allow some of steps 414 to be executed in parallel while others may execute in series. This is possible if the steps executing in parallel do not need data from the context that was inserted or added during a previous step. In the steps of FIG. 5, retrieve application state 502 and get application settings 506 do not require any information from the context and can be executed in parallel to one another. Thus is just one example of parallelism and other may exist.

[0120] By executing these steps in parallel (e.g., calling the relevant functions in parallel), the efficiency and speed of executing the entire transaction path 412 is improved. This improvement is notable in comparison to the prior monolithic designs, where steps necessarily are executed one after the other. Nonetheless, when executing steps in parallel, the data that is added to the context by such steps should be stored in separate areas of the context to avoid collisions or race conditions.

[0121] When some of steps 414 are to be executed in series, this may be due to the output of one step being needed in order to execute the next in the specified order. As an example, some of transaction paths 412 may indicate that the following steps are to be executed in series: handle request 508, generate result 510, finalize transaction 512, end transaction 514, and log transaction 516.

E. State Roll-Back and Roll-Forward

[0122] In the event that something goes wrong during the execution of a transaction path 412, orchestrator 416 can control whether transaction path 412 will halt and roll the context backward (roll-back), or retry one or more of steps 414 (roll-forward). Since steps 414 are independent, they have no knowledge of and minimal or no impact on the execution of other of steps 414 or transaction path 412 as a whole. Thus, each of steps 414 need only rely on the orchestrator 416 for invocation if execution is to be repeated.

[0123] As an example, consider steps 500-516 of FIG. 5. These steps may be executed in the order shown (e.g., session management 500, then retrieve application state 502, and so on, ending with log transaction 516). These steps may be logically divided at a boundary between generate result 510 and finalize transaction 512. If generate result 510 is successful, the boundary is crossed and any failure of the transaction processing system beyond this point for any subsequent step(s) may cause those step(s) to be retried, possibly more than once. If a failure of the transaction processing system occurs before the boundary, the processing of all previous steps may be rolled back and the transaction path 414 may be restarted. Each of transaction paths 412 may have such a boundary, possible in different locations in accordance with the ordering of steps 414 and application-level requirements.

[0124] The logic behind this approach is that if generate result 510 provides a valid result, the transaction should be completed by rolling it forward. If such a result is not obtained, the user did not receive a result and the roll-back can be used to generate a new result.

[0125] If a roll-forward cannot be executed and completed after some time, a manual completion of the transaction path 412 may be employed, or the transaction should be voided. The transaction processing system design is such that it is clear which transactions are within a transaction path 412 and thus not complete. Notably, these incomplete transactions can be determined by querying database 418 for the context of transactions and determining which have not yet completed.

F. Transaction Path Extensions

[0126] It is possible to add new steps to steps 414 and to create new instances of transaction paths 412 with new combinations of steps 414. Particularly, transaction paths 412 are designed so that the steps 414 in each can be configurable. This provides flexibility to have different transaction paths 412 for different use cases. Some transaction paths 412 may require more of steps 414 than others, and can be configured to meet the needs of a specific server application, operator, geographic location of user, and so on.

[0127] Examples of extensions to server applications include new steps to: obtain, calculate, or otherwise determine custom application settings; request and receive application state from different external services; perform new test functions; and/or replaying parts or all of a transaction without the replay impacting the context of a transaction path 412.

[0128] Examples of extensions for a specific operator include new steps 414 to: provide in-transaction communication relating to the specific operator, obtain user-personalized state from external services relating to the specific operator; and/or distribute the functionality of one or more of steps 414 to external services relating to the specific operator.

[0129] Examples of extensions relating to the geographic location of the user include new steps to: check a configuration or setting specific to a particular geographic location; steps dedicated to communicating to a geographically-specific external service; and/or steps 414 that limit transactions in a manner specific to a particular geographic location.

[0130] A new server application can be developed and can either use a previously defined transaction path 412 or a new transaction path 412 may be created. In some cases, existing

steps **414** may be reused for this new transaction path **412** or new steps **414** may be created for use with it.

G. Pre-Processing and Post-Processing

[0131] As noted above, pre-step processing **520** and/or post-step processing **522** may be used to modify the behavior of a transaction path. Pre-step processing **520** can be used to modify and/or use the context of a transaction before one of steps **414** is performed, and post-step processing **522** can be used to modify and/or use the context of a transaction after one of steps **414** is performed. In this fashion, pre-step processing **520** and/or post-step processing **522** can be used to provide exception handling, rare events, or premium functionality that is not typically part of a transaction, including the creation of novel and highly-personalized server application behavior for at least some users.

H. Load Balancing

[0132] The decomposed server architecture **400** has been designed to scale with the number of users and traffic load. It allows for load balancing of traffic across multiple server devices and across step-based logic. One or more load balancers may facilitate this process. In general, a load balancer receives traffic from a source, selects a target from multiple possible targets, and routes the traffic to the selected target.

[0133] FIG. 7 provides an example of decomposed server architecture **400** with strategically-located load balancers. Notably, load balancer **700** may be placed in front of two or more instances of orchestrator **416**. This allows the execution of different transaction paths **412** to be spread across multiple orchestrators. Likewise, load balancer **704** may be placed between the instances of orchestrator **416** and two or more instances of operator bridge **422**. This allows various steps **414** that may invoke operator bridge **422** (e.g., handle request **508** step as shown) to be spread across these instances of operator bridge **422**. Similarly, load balancer **706** may be placed between the instances of orchestrator **416** and two or more instances of external service **708**. This allows various steps **414** that may invoke external service **708** (e.g., generate result **510** step as shown) to be spread across these instances of external service **708**. These are just examples and other possible load balancer arrangements exist, e.g., for more or fewer of steps **414**.

[0134] These load balancers may employ various disciplines to determine a load-balanced target for incoming traffic. For example, round-robin (iteratively cycling through all targets), random (selecting a target based on a predetermined probability distribution), or hash-based (determining a target based on applying a hash function to part of the incoming traffic) load balancing may be used. Other disciplines may consider the size of the job to be carried out by the target, information from application metadata **408**, information from operator metadata **410**, and/or information from the context stored in database **418**.

[0135] With this load balancing, a message from any of steps **414** that needs to be transmitted to another computing device or executing process can be sent to any one of a pool of such computing devices or executing processes for service. Since these external services are stateless, it does not matter which is selected. For example, the handle request **508** step can be routed to any instance of operator bridge **422**.

[0136] Thus, each of steps **414** can use a load balancer and scale up with its traffic demand. The responsibility of performance resides with each step. The orchestrator **416** can scale up horizontally (horizontal scaling refers to increasing the capacity of a system by adding additional nodes, as opposed to increasing the capability of the existing nodes), and each step can also scale up horizontally. The result is an architecture and system that is highly scalable. Importantly, the performance of each area can be clearly monitored and identified, allowing for optimizations and performance improvements to be targeted to particular areas in need of improvement or otherwise forming a bottleneck. As example, a step of steps **414** that is the slowest in comparison to the others of steps **414** will be a candidate for performance improvement.

[0137] Additionally, once such a step is using a load balancer, it allows for adaptive scaling. A threshold can be set for the resource usage or processing time for a step. If this threshold is reached, then the load balancer technology can cause additional service capacity (e.g., more physical server devices, virtual servers, and/or software processes) to be added to the pool. This facilitates sharing the traffic across more services so that the desired performance is more likely to be achieved by the step. If the traffic volume reduces, the additional services can be scaled down in a converse manner. As a result, automatically scaling up and down the performance of each step is possible.

[0138] It is also possible for an instance of an external service behind a load balancer to fail. In a simple example, the external service can crash and become unavailable. In this situation, the load balancer may be able to detect the failure (e.g., due to lack of a response to a message) and at least temporarily remove the failed instance from its load balancing targets. Orchestrator **416** can invoke the step again, and then the load balancer may direct the message to an instance of the external service that is still operational. This, effectively bypasses the failure and results in the continuation of the transaction path. The instance that failed can be restored back to health in its own time (and added back to the targets of the load balancer) without blocking any execution.

[0139] It has been a long-standing issue in multiple industries that when a new server application goes live, there can be huge volumes of user attempting to access and use the server application on the first few days after its release. This increase in traffic can be unpredictable, as server application may be more or less successful than expected. The result is that there may be an increase in latency of server application responses, as well as unresponsiveness to users.

[0140] Thus, in days soon after launch, there may be a need to seamlessly scale the transaction processing system without impacting any other server applications. Server applications should have consistent latency and performance for the users, irrespective of the volume of traffic. When designed as described herein, the transaction processing system can be scaled up and down with demand for any given transaction path or step. Such automatic scaling allows an improved and seamless user experience with reduced variability of latency.

I. Improved Loading Time

[0141] In reference to FIG. 4, loading time is defined from the point that a user selects a server application by way of operator web site **402** until the user can interact with this

server application. When a server application is loading, the system carries out at least two tasks. First, it may verify that the user has permission to access the server application (e.g., by checking with one or more of location/geoblocking service **406**, application metadata **408**, and/or operator metadata **410**). Second, it may retrieve any active state for this server application that is either stored in database **418** or elsewhere. This may not be the first time the user is accessing the server application and, as such, it may retrieve the state as of the time the user last accessed the server application. Obtaining this state is referred to herein as a “refresh” request. Such state may, for example, cause client device **300** to display a previous result relating to execution of the server application.

[0142] A long-standing problem is to reduce this loading time so that the user only waits as much time as needed. The longer the user waits, the more likely they are to exit client application **300A**. Decomposed server architecture **400** addresses this issue.

[0143] When the user initially launches the server application, the operator can supply an authentication token as part of the URL of application launcher **404**. This allows the transaction processing system to perform a server-to-server callback to the operator in order to validate that the user is permitted to launch and use this particular server application.

[0144] Application launcher **404** can invoke the session management **500** step to create a session for the user, and provide a unique session ID to operator web site **402**. Application launcher **404** can then look up the application address (e.g., another URL) to use for this operator and server application from operator metadata **410**. Application launcher **404** can then provide the application address to client application **300A** so that client application **300A** can start loading the server application content.

[0145] The callback to the operator is made to validate the player, but importantly, the server application can continue to load while the operator is responding. Once the server application is sufficiently loaded enough, it can perform a refresh request. The loading is only slowed down while checking, based on the authentication token, that the user is allowed to access the server application. Since the transaction path is responding to a refresh request, it can be optimized to respond quickly to the request. Parallelism achieved in this initial phase of loading a server application and authenticating with an operator is possible through new features provided by decomposed server architecture **400**.

J. Security

[0146] Since the program logic of any of steps **414** can be outside of orchestrator **416**, the program logic of such a step can also reside in a different virtual private cloud (VPC) from orchestrator **416**. A VPC is a secure, isolated private cloud environment that can be hosted within a public cloud environment. It enables allocation and use of resources in a virtual network, including the selection of IP address ranges, creation of subnets, and configuration of routing tables and network gateways. A VPC can provide advanced security features such as security groups and network access control lists, which offer granular control over inbound and outbound traffic to and from individual resources. Additionally, VPCs support various virtualized and software-based

resources, such as virtual machines, containers, and databases, and allow for scaling of resources according to demand.

[0147] Use of VPCs allows for greater control over the security of communications between steps **414** and external services. In the example of FIG. 4, operator bridge **422** and transaction database **426** can reside in different VPCs (perhaps the same VPCs as the steps by which they are respectively invoked). Thus, a new server application can be deployed without needing to modify the VPC where orchestrator **416** is disposed.

K. Reduction in Hardware Use

[0148] The hardware used in decomposed server architecture **400** can execute on cloud-based and/or virtualized technologies, for example Kubernetes. This can improve efficiencies over existing systems. Particularly, the computational cost for the hardware and software resources used in determining a result can be accurately measured. Currently, no system known is able to measure performance and resource costs this precisely.

[0149] The underlying resources used with any transaction path can be switched out for more cost-effective resources as needed. This is possible due to the flexibility and scalability of the system. Since the system can be scaled horizontally, there is a choice as to how many instances of orchestrator **416** and/or any of steps **414** are executing simultaneously. There could be a few instances running on high-power hardware resources (e.g., with faster processors and more memory), or more instances executing on lower-power hardware resources (e.g., with slower processors, and/or less memory).

[0150] Decomposed server architecture **400** is also highly flexible in terms of programming languages and database technologies used. This polyglot arrangement also allows for optimizations to reduce computational cost. For example, the choice of database technology can easily factor into the cost of transactions or steps. This flexibility is something not achieved with traditional (known) architectures.

L. Transaction Logging

[0151] Often the final step in a transaction path **412**, data from the transaction and context may be logged by log transaction **516** step to transaction database **426** (which could be based on, for example, Apache Kafka). Once this step has logged the data, it orchestrator **416** may provide the response to the client. This final step makes logging a high priority of a transaction path **412**. Usually this is not the case and logging occurs separately, with a low priority, or not at all.

[0152] This logging of transaction data it allows other systems to receive the data in either near-real time or in batches. For server applications that require real-time monitoring (e.g., for correctness, stability, or performance purposes), they can receive the data needed and take action in an expedited fashion.

IV. Example Decomposed Web-Based Service

[0153] FIG. 8 depicts a decomposed web-based service architecture **800** in line with the discussion above. Service architecture **800** could facilitate, for example, a storage services, a customer relationship management service, a

videoconferencing service, a messaging service, a blogging service, or some other form of service. This is just one example of an application that can employ server architecture 400 for improved performance.

[0154] Thus, service architecture 800 may include some or all features of server architecture 400. Though not shown in FIG. 8, these can include software modules for operator web site 402, application launcher 404, location/geoblocking service 406, application metadata 408, operator metadata 410, and orchestrator 416.

[0155] Steps 414 of a transaction path 412 may include steps 500-516 as described above. Here, transaction path 412 indicates that get application settings 506 can be performed in parallel to retrieve application state 502 and establish transaction ID 504. All other of steps 500-516 are performed serially as shown by the arrows.

[0156] Session management 500 may interact with operator services 424 by way of operator bridge 422. For example, session management 500 may authenticate the account of a user attempting to access the web-based service. Similarly, handle request 508 may interact with operator services 424 by way of operator bridge 422 to provide representations of user requests to operator services 424, generate result 510 may interact with operator services 424 by way of operator bridge 422 to receive the results of user requests from operator services 424, and finalize transaction 512 may interact with operator services 424 by way of operator bridge 422 to confirm that the user has accepted these results.

[0157] As noted above, log transaction 516 may provide data, to transaction database 426, representing each transaction. Monitoring services 428 may obtain this data either in real-time or as needed.

[0158] Service architecture 800 can take advantage of any of the features described above. As shown, it uses parallelism. It can employ roll-back and roll-forward features, e.g., with a boundary set between generate result 510 and finalize transaction 512. Various pre-processing and post-processing steps can be inserted before, after, and/or between the steps shown. Further, load balancing and loading time improvements are possible. Other advantages are also possible.

V. Example Decomposed Entertainment Service

[0159] Online entertainment services have grown in popularity in recent years, as more of the world's population now has reliable, high-speed Internet access. Such entertainment services may include streaming audio, streaming video, e-books, audio books, various forms of social media, and so on. Online gaming is also a popular and growing set of entertainment services with genres, styles, and platforms that can vary dramatically (e.g., app-based, web-based, platform-specific, etc.).

[0160] Some specific types of online games suffer from many of the drawbacks of traditional transaction processing systems and can benefit from the decomposed server architecture described herein. Particularly, wagering games can be arranged around specific transaction paths that interact with external services in various ways—thus, they stand to benefit from the technical improvements of the decomposed server architecture. Non-exclusive examples of such wagering games include reel-based games, card games, dice games, and various types of mechanical games.

[0161] Reel-based games are primarily slot games that consist of spinning reels with symbols. Users win when the

reels stop spinning to display specific combinations of symbols along a payline. Examples include 3-reel and 5-reel variations, and symbols can be fruits, bars, numbers, or other types such as those related to a specific theme. Some real-based games include multiple paylines, free spins, bonus rounds, and/or wild symbols, each of which increases the chance of winning.

[0162] Card games involve wagering on the outcome of hands dealt from a deck of cards (e.g., a standard deck of 52 cards). These games include blackjack, baccarat, and poker as some possible examples. In blackjack, users play against a dealer (which could be virtual) and aim to get a hand total as close to 21 as possible without going over. In baccarat, users bet on which of two hands (the “user” or the “house”) will have a hand total closest to 9. In poker, the user attempts to achieve a hand of cards that is higher-ranked than any of their opponents. There are many variations available including Texas Hold'em (users receive two private cards and try to make the best hand with five community cards) and five-card draw (users are dealt five cards and can exchange some or all of them for new ones in hopes of making a better hand).

[0163] Other types of games are based on (virtual) mechanics with randomized outcomes, such as spins of specific devices and/or die rolls. These include roulette, craps, and bingo, for example. In roulette, users wager on where a ball will land on a spinning wheel. Locations on the wheel include specific numbers, colors (e.g., red or black), odd or even numbers, and various combinations thereof. In craps, users wager on the outcomes of a die roll, or a series of die rolls, of a pair of dice. In bingo, users match numbers printed in different arrangements on cards with the numbers drawn at random, with the users seeking to obtain certain patterns of selected numbers as arranged on their tiles.

[0164] Regardless of type, each of these games may include a set of steps that a server device may perform on behalf of a user including: user authentication, establishing a new game, allowing a user to place a wager, generating a result, optionally modifying the user's account profile based on the result, and so on. Some of these steps may involve accessing external services. For example, user authentication, placement of a wager, and/or modifying the user's profile may require that a server device interact with an external service in order for the respective step to be completed.

[0165] Thus, gaming applications can be arranged as forms of transaction paths (“paths”) as described above. In other words, such a game may define one or more transaction paths, each transaction path including some number of steps. Accordingly, many of the aforementioned advantages of the decomposed server architecture described above apply these gaming applications.

A. Additional Drawbacks of Monolithic Server Architectures

[0166] In addition to the drawbacks of monolithic and/or inflexible server architectures described above, gaming applications also suffer from further drawbacks when implemented based on traditional architectures. For example, establishment of a new transaction (e.g., turn of a game), or continuation of a previous transaction that was paused or exhibited a failure, can involve a varied and growing number of operations.

[0167] For example, regulated markets are increasingly adding more controls around the wager values that users are allowed to place. Server-based logic now needs to be aware of the country or region the user is playing in and the latest configuration for a wager setting. In addition, operators are also increasingly controlling the maximum amount that a user can wager. Thus, there are multiple levels of wager settings that may be checked at or near the beginning of a transaction. These levels include geographic region, operator, and the game-specific settings. Each level may have a different priority, and the combination of these levels of settings determines whether and how the user can interact with the gaming application.

[0168] Other regulated market requirements may limit which features in a gaming application are available to the player. For example, in some geographic locations, a gaming application may not be allowed to offer in-game feature purchases or autoplay functionality. These are additional examples where a step needs to determine whether a feature is allowed in the relevant geographic location, and then block the transaction if required.

[0169] Other geographic locations and/or operators allow or require that a user limit wagers (e.g., per transaction wagers or total wagers over a period of time or a set of transactions) and/or session times (e.g., the amount of time that the user has been playing the gaming application continuously or within a fixed period such as a day). These limits are very complex to track, each requiring additional state to be stored and checked during a transaction.

B. Example Decomposed Architecture

[0170] FIG. 9 depicts a decomposed gaming service architecture 900 in line with the discussion above. Service architecture 900 is similar to that of service architecture 800, but includes a different arrangement of steps into a different transaction path. This transaction path could be employed to facilitate various types of online games, such as reel-based games, card games, and other types of wagering games. Notably, FIG. 9 depicts just one possible transaction path for gaming applications and others may exist.

[0171] Additionally, service architecture 900 may include some or all features of server architecture 400. Though not shown in FIG. 9, these can include software modules for operator web site 402, application launcher 404, location/geoblocking service 406, application metadata 408, operator metadata 410, and orchestrator 416.

[0172] In the context of service architecture 900, each operator may provide and/or manage the execution of one or more gaming applications. In some cases, different operators may provide the same or similar gaming applications with minor or cosmetic variances (e.g., different skins, wagering constraints, etc.).

[0173] As was the case for service architecture 800, steps 414 of a transaction path 412 may include steps 500-516 as well as other steps described above. Again, transaction path 412 indicates that get application settings 506 can be performed in parallel to retrieve application state 502 and establish transaction ID 504. All other steps are performed serially as shown by the arrows. Nonetheless, various embodiments might involve different orderings, some of which are discussed below.

[0174] Session management 500 may interact with operator services 424 by way of operator bridge 422 (which, as noted may be a gateway, proxy, or interface to operator

services 424). For example, session management 500 may authenticate the account of a user attempting to access the gaming application.

[0175] Retrieve application state 502 may involve obtaining (e.g., from database 418 or elsewhere) the application state and/or the user's progress within the gaming application as represented by steps of the transaction path. For example, retrieve application state 502 may add this application state to the transaction's context.

[0176] Establish transaction ID 504 may involve generating a transaction identifier (e.g., a unique number or token) representing the current transaction.

[0177] Get application settings 506 may involve obtaining operational settings for the gaming application (e.g., from application metadata 408, operator metadata 410, and/or user-specific information). This may involve wagering-related setting such as limits on various types or patterns of wagers. Get application settings 506 may add these settings to the transaction's context.

[0178] Pre-step processing 520 may involve any logic to execute and/or external services to invoke before a result is generated. Thus, these operations may be specific to the particular gaming application and/or operator involved in the transaction. The outcome of pre-step processing 520 may be added to the transaction's context.

[0179] Generate game result 510 may involve determining a result (e.g., an outcome of a game or a turn of a game). This result may be determined locally by the gaming application (as shown) or in conjunction with the operator by way of operator bridge 422 and operator service 424. Generate game result 510 is a renaming of the generate result 510 step to signify that it may entail gaming-specific processing and results. A representation of the generated game result may be added to the transaction's context.

[0180] Place wager 508 may involve receiving and processing a wager from client application 300A, as well as possibly providing a representation of the wager to the operator by way of operator bridge 422 and operator service 424 for approval. Place wager 508 is a renaming of the handle request 508 step to signify that it may entail gaming-specific processing and results. A representation of the wager may be added to the transaction's context.

[0181] Notably, the ordering of steps 508 and 510 are shown differently in FIG. 9 than they were in FIG. 8. In gaming applications provided by service architecture 900, either order may be employed. In the ordering of FIG. 9, the game result produced by generate game result 510 may be hidden or withheld from the user until the user's wager of place wager 508 is received and approved (e.g., by local policy of the gaming application or by operator services 424). In some cases, either ordering of steps 508 and 510 may involve some indication of progress being shown by way of client application 300A from: (i) when the wager and/or its approval is received, until (ii) the result is available to the gaming application or at some point thereafter. For instance, an animation such as the spinning of reels may be displayed on client device 300A to hide any latency involved with generating the result.

[0182] Post-step processing 522 may involve any logic to execute and/or external services to invoke after a result is generated. Thus, these operations may be specific to the particular gaming application and/or operator involved in the transaction. The outcome of pre-step processing 520 may be added to the transaction's context.

[0183] In FIG. 9, finalize transaction 512 has been broken up into two steps: save final state 512A and update user profile 512B. Save final state 512A may involve logic to record the final game state (e.g., a win or loss, a number of points, a payoff of the wager, a log of the transaction) to the transaction's context and/or other locations. Update user profile 512B may involve updating operator service 424 by way of operator bridge 422 regarding any credits or debits to be made to the user's profile with the operator. For example, a wager with no associated payout from generate game result 510 (or a payout less than the value of the wager) may cause the user's profile to be debited, whereas a wager with an associated payout from generate game result 510 of at least the value of the wager may cause the user's profile to be credited.

[0184] End transaction 514 may involve completing and ending the transaction. For example, this may involve deallocating memory used for the transaction and otherwise marking the transaction as complete, as well as possibly notifying the operator that the transaction is complete by way of operator bridge 422 and operator service 424.

[0185] As was the case in FIG. 8, log transaction 516 may provide data, to transaction database 426, representing each transaction. Monitoring services 428 may obtain this data either in real-time or as needed.

[0186] Service architecture 900 can take advantage of any of the features described above. As shown, it uses parallelism. It can employ roll-back and roll-forward features, e.g., with a boundary set between generate result 510 and finalize transaction 512. Various pre-processing and post-processing steps can be inserted before, after, and/or between the steps shown. Further, load balancing and loading time improvements are possible. Other advantages are also possible.

C. Extensions to Transactions

[0187] Transaction paths 412 may be designed so that their steps 414 can be configurable. This provides flexibility to have different transaction paths 412 for different use cases, reasons, and/or events. Some transaction paths 412 may require more steps 414 than others, and can be configured to meet the needs of a specific gaming application, product, operators, geographical location, or the like. Some examples follow.

[0188] For a transaction path 412 designed for a particular geographical location, additional steps 414 may: check a new configuration specific to particular geographical location, communicate to an external regulatory system of the particular geographical location, and/or calculate and enforce a new limitation (e.g., bet or session time limits, feature usage, etc.).

[0189] For a transaction path 412 designed for a specific gaming application or gaming application type, additional steps 414 may: determine specific wager settings and configurations that are custom to the specific gaming application or gaming application type, retrieve state from other systems which can then feed into generate game result 510, perform custom tests with test data for the specific gaming application or gaming application type, and/or formulate a transaction path that can replay a gaming application result without affecting the user's payout or balance.

[0190] For a new or existing transaction path 412 designed to incorporate new steps 414, these steps may include additive features like bonus games and jackpots, external configuration, settings or data (e.g., loyalty points or a

power-up functionality), and/or provide a deep integration into another system (e.g., a multiplayer live dealer system could execute user wagers in a particular transaction path 412 and have custom steps 414 to obtain the live dealer results from an external system).

[0191] A feature, gaming application, or gaming application type can be developed and can either use a current transaction path 412, or can create a new transaction path 412 and re-use existing steps 414. This reuse of steps 414 and customizable transaction paths 412 enables quick development, testing, and deployment of new applications.

D. Extensions to Steps

[0192] As noted, pre-step processing 520 and post-step processing 522 may be placed in various locations of a transaction path. As shown in FIG. 9, pre-step processing 520 can be performed before generate game result 510, and may use and/or modify the context before such a result is generated. Also, post-step processing 522 can be performed after place wager 508, and may use and/or modify the context at that point. Other possibilities exist.

[0193] Either or both of pre-step processing 520 and post-step processing 522 may be used to control unusual, bonus, or unexpected functionality that is not typically part of a transaction path for a gaming application. This may allow the creation of novel and highly personalized experiences for users.

[0194] Pre-step processing 520 allows the use of mechanics that can check the context for any active influence on an upcoming or next execution of generate game result 510. The active influence could be, for example, that a user is engaging in a threshold extent of the transaction path (e.g., placing the user's 100th or 1000th wager in the gaming application). In return, the user may be granted a certain number of free transactions with the gaming application (e.g., free games) or access to a different bonus transaction path (e.g., a bonus game).

[0195] These types of features can be controlled per operator, such that each operator may include an indication of their threshold in the context, and post-step processing 522 may increment a counter of transactions per user. When pre-step processing 520 determines that a user's counter has met or exceeded the threshold, it may also update the context to indicate that the awarded free transactions or access to the different bonus transaction path are available to the user.

[0196] Another example pre-step processing 520 is determining (e.g., from the context) that a user is eligible for different type or additional result (e.g., an increase on returns or a reward for loyalty). This might be because the user owns an NFT or another type of digital token. This information can be included in the context and provided to generate game result 510.

[0197] Post-step processing 522 allows for changes or variations to be processed after generate game result 510 completes. Post-step processing 522 may include, for example, an increase to the user's return from the game result which is then stored in the context. In some cases, this might only apply to a certain type of result or after a threshold number of results have been generated for the user.

[0198] For avoidance of doubt, any one or more (or all) of these changes made based on pre-step processing 520 and post-step processing 522 can be recorded in the context. Thus, an audit log could be generated for the changes based on the context.

E. Certification

[0199] Some geographic locations are subject to certification of transaction paths and their associated applications. Such certification procedures may be voluntary or mandated by law or regulation. Regardless, certification can be a complex process involving manual review of the underlying program code of a gaming application in order to determine its behavior and properties. Moreover, each time such an application changes (e.g., a new version is ready for deployment), the program code may require recertification before it can be made available to users. Since traditional transaction processing systems are monolithic, this means that any change to such a system would result in the entire system being recertified at great effort and cost.

[0200] In service architecture 900, steps 414 are independent (or at least minimally related to one another). This allows for steps 414 to be changed, removed, or added without the need for the entire application to be recertified. In these cases, the certification effort can focus only on the blocks of program code that changed and/or those that are considered critical for application functionality.

[0201] For example, traditional certification requires the hashing of critical files (e.g., by way of a hash function such as SHA-256 or some other type of one-way function) and then determining when these file have changed by comparing hashes of new versions of files to hashes of their respective older versions. When corresponding hashes are different, then the content of the file (e.g., its program code) has changed and may require recertification.

[0202] Since the steps 414 can be broken down into separate files in service architecture 900, this hashing method is faster and easier to track and manage. For instance, modular files that were not changed generally do not require review. Moreover, the file size will be reduced compared to what is currently the case for existing systems. In this manner, new transaction paths 412 and steps 414 can be deployed without the need for all steps 414 to be recertified.

VI. Example Operations

[0203] FIGS. 10 and 11 are flow charts illustrating example embodiments. The processes illustrated by FIGS. 10 and 11 may be carried out by a computing device, such as computing device 100, and/or a cluster of computing devices, such as server cluster 200. However, the processes can be carried out by other types of devices or device subsystems. For example, the processes could be carried out by a portable computer, such as a laptop or a tablet device.

[0204] The embodiments of FIGS. 10 and 11 may be simplified by the removal of any one or more of the features shown therein. Further, these embodiments may be combined with features, aspects, and/or implementations of any of the previous figures or otherwise described herein.

[0205] Block 1000 of FIG. 10 may involve receiving, by a web site, selection of a server application, wherein the server application is associated with a path, wherein the path includes a set of steps, and wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions. Here, such a discrete set of functions are generally well-defined to perform specific tasks, may execute independently from one another, and may be implemented in separate blocks or units of program code.

[0206] Block 1002 may involve receiving, by an application launcher, a request on behalf of a client device. This request may be from the client device by way of the web server. The request may also be from the same client device that selected the server application. For instance, the selection of the server device may cause the request to be provided to the application launcher.

[0207] Block 1004 may involve determining, by the application launcher, the server application based on the request, and also based on metadata that relate to the server application and an operator of the web site.

[0208] Block 1006 may involve providing, by the application launcher, a client application to the client device, wherein the client application is configured to interact with the server application.

[0209] Block 1008 may involve executing, by an orchestrator, the path, wherein at least some of the set of steps include interaction with one or more external services, and wherein a context of the execution is written in a database between execution of at least some of the set of steps.

[0210] In some examples, the web site is also configured to allow selection of a second server application, wherein the second server application is associated with a second path, wherein the second path includes a second set of steps, wherein each of the second set of steps is respectively associated with second program logic to perform a second discrete set of functions, wherein the metadata also relates to the second server application, wherein the application launcher is also configured to: receive a second request from the client device, determine the second server application based on the second request and the metadata, and provide a second client application to the client device, wherein the second client application is configured to interact with the second server application, wherein the orchestrator configured to execute the second path, wherein at least some of the second set of steps include interaction with the one or more external services, and wherein a second context of the execution is written in the database between execution of at least some of the second set of steps.

[0211] In some examples, the software components further comprise a second web site configured to allow selection of a second server application, wherein the second server application is associated with a second path, wherein the second path includes a second set of steps, and wherein each of the second set of steps is respectively associated with second program logic to second perform a second discrete set of functions, wherein the metadata also relates to the second server application and a second operator of the second web site, wherein the application launcher is also configured to: receive a second request from the client device, determine the second server application based on the second request and the metadata, and provide a second client application to the client device, wherein the second client application is configured to interact with the second server application, wherein the orchestrator configured to execute the second path, wherein at least some of the second set of steps include interaction with the one or more external services, and wherein a second context of the execution is written in the database between execution of at least some of the second set of steps.

[0212] In some examples, providing the client application to the client device is based on determining, by way of a geolocation service, that the client device is permitted to access the server application.

[0213] In some examples, the client application is executable within a web browser of the client device.

[0214] In some examples, the set of steps include one or more of: session management, retrieving application state, establishing a path identifier, getting application settings, handling requests from the client device, generating results, finalizing execution of the path, ending execution of the path, or logging data relating to execution of the path.

[0215] In some examples, the path indicates that: at least two of the set of steps can be performed in parallel to one another, and a pair of steps in the set of steps are to be performed serially.

[0216] In some examples, a particular step of the set of steps has either pre-step processing or post-step processing defined thereof, wherein the pre-step processing occurs before execution of the particular step and is configured to modify the context, and wherein the post-step processing occurs after execution of the particular step and also is configured to modify the context.

[0217] In some examples, each of the set of steps is configured to execute independently of other steps of the set of steps.

[0218] In some examples, the path defines a boundary separating the set of steps into a first group of steps and a second group of steps, wherein the first group of steps are configured to be performed before the second group of steps, wherein a failure that occurs during execution of the first group of steps causes the context to be rolled back and the path to be restarted, and wherein a further failure that occurs during execution of the second group of steps causes at least some of the second group of steps to be retried.

[0219] Some examples may also include: a plurality of orchestrators, each configured to execute paths, wherein the orchestrator is one of the plurality of orchestrators; and a load balancer configured to distribute representations of the paths across the plurality of orchestrators according to a load balancing discipline.

[0220] In some examples, the plurality of orchestrators scales in number proportionally to orchestrator demand.

[0221] In some examples, a particular step of the set of steps is configured to interact with a particular external service of the one or more external services, wherein a plurality of server components are configured to perform the particular external service. These examples may further include a load balancer configured to distribute interactions across the plurality of server components according to a load balancing discipline.

[0222] In some examples, the plurality of server components scales in number proportionally to server component demand.

[0223] In some examples, the application launcher is configured to facilitate loading of the server application in parallel to determining whether a user associated with the client device is permitted to access the server application.

[0224] In some examples, the orchestrator and the one or more external services are hosted within different virtual private cloud systems.

[0225] In some examples, the context represents an intermediate or final state of the execution of the path.

[0226] In some examples, the orchestrator receives a representation of the path from the application launcher.

[0227] In some examples, at least some of the external services operate in a stateless fashion.

[0228] The features, alternative examples, and additional examples discussed in the context of FIG. 10 may be arranged in various combinations. In other words, any one or more of these alternative examples or additional examples can be selected for inclusion with blocks 1000, 1002, 1004, 1006, and 1008.

[0229] Block 1100 of FIG. 11 may involve receiving, by an application launcher, a message on behalf of a client device.

[0230] Block 1102 may involve determining a server application based on the message and metadata relating to the server application, an operator of the server application, or a user of the server application, wherein the server application is associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions, wherein a generate result step of the set of steps causes the server application to determine an outcome of executing the path, and wherein a handle request step of the set of steps causes the server application to provide, to an external service, a representation of a request related to the outcome.

[0231] Block 1104 may involve providing, by the application launcher, a client application to the client device, wherein the client application is configured to interact with the server application.

[0232] Block 1106 may involve executing, by an orchestrator, the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.

[0233] In some examples, the software components further comprise a web site configured to allow selection of the server application from a menu of server applications associated with the operator or with one or more other operators.

[0234] In some examples, the path specifies that pre-step processing is to be performed prior to the generate result step, and wherein the pre-step processing comprises: reading, from the database, at least part of the context; determining, from the context as read, that an additional feature applies to the generate result step; and performing the generate result step in accordance with the additional feature.

[0235] In some examples, determining that the additional feature applies to the generate result step comprises determining that the context indicates that the user has engaged with the server application a threshold number of times.

[0236] In some examples, the path specifies that post-step processing is to be performed after the handle request step, and wherein the post-step processing comprises: reading, from the database, at least part of the context; determining, from the context as read, that an additional modification of the context is to be applied after the handle request step; and writing, to the context, an update based on the additional modification.

[0237] In some examples, the path specifies that pre-step processing is to be performed prior to the generate result step, and wherein the pre-step processing comprises: reading, from the database, at least part of the context; determining, from the context as read, that a constraint applies to the handle request step; and performing the handle request step in accordance with the constraint.

[0238] In some examples, the constraint relates to an amount specified by the user as part of the handle request step.

[0239] In some examples, the server application is a gaming application, and wherein the outcome is based on an iteration of the gaming application.

[0240] In some examples, providing the client application to the client device is based on determining, by way of a geolocation service, that the client device is permitted to access the server application.

[0241] In some examples, the client application is executable within a web browser of the client device.

[0242] In some examples, the path indicates that: at least two of the set of steps can be performed in parallel to one another, and a pair of steps in the set of steps are to be performed serially.

[0243] In some examples, the path defines a boundary separating the set of steps into a first group of steps and a second group of steps, wherein the first group of steps are configured to be performed before the second group of steps, wherein a failure that occurs during execution of the first group of steps causes the context to be rolled back and the path to be restarted, and wherein a further failure that occurs during execution of the second group of steps causes at least some of the second group of steps to be retried.

[0244] In some examples, the boundary is between the handle request step and a further step that immediately follows either: the handle request step, or any post-processing step related to the handle request step.

[0245] Some examples may further involve: a plurality of orchestrators, each configured to execute paths, wherein the orchestrator is one of the plurality of orchestrators; and a load balancer configured to distribute representations of the paths across the plurality of orchestrators according to a load balancing discipline.

[0246] In some examples, a particular step of the set of steps is configured to interact with a particular external service of one or more external services, wherein a plurality of server components are configured to perform the particular external service. These examples may further involve a load balancer configured to distribute interactions across the plurality of server components according to a load balancing discipline.

[0247] In some examples, the application launcher is configured to facilitate loading of the server application in parallel to determining whether a user associated with the client device is permitted to access the server application.

[0248] In some examples, the orchestrator and one or more external services invoked by the path are hosted within different virtual private cloud systems.

[0249] In some examples, the context represents an intermediate or final state of the execution of the path.

[0250] In some examples, the orchestrator receives a representation of the path from the application launcher.

VII. Closing

[0251] The present disclosure is not to be limited in terms of the particular embodiments described in this application, which are intended as illustrations of various aspects. Many modifications and variations can be made without departing from its scope, as will be apparent to those skilled in the art. Functionally equivalent methods and apparatuses within the scope of the disclosure, in addition to those described herein, will be apparent to those skilled in the art from the foregoing descriptions. Such modifications and variations are intended to fall within the scope of the appended claims.

[0252] The above detailed description describes various features and operations of the disclosed systems, devices, and methods with reference to the accompanying figures. The example embodiments described herein and in the figures are not meant to be limiting. Other embodiments can be utilized, and other changes can be made, without departing from the scope of the subject matter presented herein. It will be readily understood that the aspects of the present disclosure, as generally described herein, and illustrated in the figures, can be arranged, substituted, combined, separated, and designed in a wide variety of different configurations.

[0253] With respect to any or all of the message flow diagrams, scenarios, and flow charts in the figures and as discussed herein, each step, block, and/or communication can represent a processing of information and/or a transmission of information in accordance with example embodiments. Alternative embodiments are included within the scope of these example embodiments. In these alternative embodiments, for example, operations described as steps, blocks, transmissions, communications, requests, responses, and/or messages can be executed out of order from that shown or discussed, including substantially concurrently or in reverse order, depending on the functionality involved. Further, more or fewer blocks and/or operations can be used with any of the message flow diagrams, scenarios, and flow charts discussed herein, and these message flow diagrams, scenarios, and flow charts can be combined with one another, in part or in whole.

[0254] A step or block that represents a processing of information can correspond to circuitry that can be configured to perform the specific logical functions of a herein-described method or technique. Alternatively or additionally, a step or block that represents a processing of information can correspond to a module, a segment, or a portion of program code (including related data). The program code can include one or more instructions executable by a processor for implementing specific logical operations or actions in the method or technique. The program code and/or related data can be stored on any type of computer readable medium such as a storage device including RAM, a disk drive, a solid-state drive, or another storage medium.

[0255] The computer readable medium can also include non-transitory computer readable media such as non-transitory computer readable media that store data for short periods of time like register memory and processor cache. The non-transitory computer readable media can further include non-transitory computer readable media that store program code and/or data for longer periods of time. Thus, the non-transitory computer readable media may include secondary or persistent long-term storage, like ROM, optical or magnetic disks, solid-state drives, or compact disc read only memory (CD-ROM), for example. The non-transitory computer readable media can also be any other volatile or non-volatile storage systems. A non-transitory computer readable medium can be considered a computer readable storage medium, for example, or a tangible storage device.

[0256] Moreover, a step or block that represents one or more information transmissions can correspond to information transmissions between software and/or hardware modules in the same physical device. However, other information transmissions can be between software modules and/or hardware modules in different physical devices.

[0257] The particular arrangements shown in the figures should not be viewed as limiting. It should be understood that other embodiments could include more or less of each element shown in a given figure. Further, some of the illustrated elements can be combined or omitted. Yet further, an example embodiment can include elements that are not illustrated in the figures.

[0258] While various aspects and embodiments have been disclosed herein, other aspects and embodiments will be apparent to those skilled in the art. The various aspects and embodiments disclosed herein are for purpose of illustration and are not intended to be limiting, with the true scope being indicated by the following claims.

What is claimed is:

1. A computer-implemented method comprising:
 - determining a server application based on a client device or metadata relating to the server application, wherein the server application is associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions;
 - based on the server application, providing a client application to the client device, wherein the client application is configured to interact with the server application; and
 - executing the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.
2. The computer-implemented method of claim 1, wherein a generate result step of the set of steps causes the server application to determine an outcome of executing the path.
3. The computer-implemented method of claim 2, wherein the path specifies that pre-step processing is to be performed prior to the generate result step, and wherein the pre-step processing comprises:
 - reading, from the database, at least part of the context;
 - determining, from the context as read, that an additional feature applies to the generate result step; and
 - performing the generate result step in accordance with the additional feature.
4. The computer-implemented method of claim 3, wherein determining that the additional feature applies to the generate result step comprises:
 - determining that the context indicates that a user has engaged with the server application a threshold number of times.
5. The computer-implemented method of claim 2, wherein a handle request step of the set of steps causes the server application to provide, to an external service, a representation of a request related to the outcome.
6. The computer-implemented method of claim 5, wherein the path specifies that post-step processing is to be performed after the handle request step, and wherein the post-step processing comprises:
 - reading, from the database, at least part of the context;
 - determining, from the context as read, that an additional modification of the context is to be applied after the handle request step; and
 - writing, to the context, an update based on the additional modification.
7. The computer-implemented method of claim 1, wherein the server application is determined based at least

on the metadata, and wherein the metadata relates to an operator of the server application or a user of the server application.

8. The computer-implemented method of claim 1, wherein the server application is determined by way of selection, on a web site from a menu of server applications associated with an operator of the server application or with one or more other operators.

9. The computer-implemented method of claim 1, wherein providing the client application to the client device is based on determining, by way of a geolocation service, that the client device is permitted to access the server application.

10. The computer-implemented method of claim 1, wherein the client application is executable within a web browser of the client device.

11. The computer-implemented method of claim 1, wherein the path indicates that: at least two of the set of steps can be performed in parallel to one another, and a pair of steps in the set of steps are to be performed serially.

12. The computer-implemented method of claim 1, wherein the path defines a boundary separating the set of steps into a first group of steps and a second group of steps, wherein the first group of steps are configured to be performed before the second group of steps, wherein a failure that occurs during execution of the first group of steps causes the context to be rolled back and the path to be restarted, and wherein a further failure that occurs during execution of the second group of steps causes at least some of the second group of steps to be retried.

13. The computer-implemented method of claim 12, wherein a handle request step of the set of steps causes the server application to provide, to an external service, a representation of a request, and wherein the boundary is between a handle request step and a further step that immediately follows either: the handle request step, or any post-processing step related to the handle request step.

14. The computer-implemented method of claim 1, wherein a load balancer is configured to distribute representations of a plurality of paths across a plurality of orchestrators according to a load balancing discipline, and wherein the path is one of the plurality of paths.

15. The computer-implemented method of claim 1, wherein a particular step of the set of steps is configured to interact with a particular external service of one or more external services, wherein a plurality of server components are configured to perform the particular external service, and wherein a load balancer configured to distribute interactions across the plurality of server components according to a load balancing discipline.

16. The computer-implemented method of claim 1, wherein an application launcher is configured to facilitate loading of the server application in parallel to determining whether a user associated with the client device is permitted to access the server application.

17. The computer-implemented method of claim 1, wherein an orchestrator that executes the path and one or more external services invoked by the path are hosted within different virtual private cloud systems.

18. The computer-implemented method of claim 1, wherein the context represents an intermediate or final state of the execution of the path.

19. A non-transitory computer-readable medium storing program instructions that, when executed by one or more

processors of a computing system, cause the computing system to perform operations comprising:

- determining a server application based on a client device or metadata relating to the server application, wherein the server application is associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions;

- based on the server application, providing a client application to the client device, wherein the client application is configured to interact with the server application; and

- executing the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.

20. A computing system comprising:

- one or more processors;

- memory; and

program instructions, stored in the memory, that upon execution by the one or more processors cause the computing system to perform operations comprising:

- determining a server application based on a client device or metadata relating to the server application, wherein the server application is associated with a path, wherein the path includes a set of steps, wherein each of the set of steps is respectively associated with program logic to perform a discrete set of functions;

- based on the server application, providing a client application to the client device, wherein the client application is configured to interact with the server application; and

- executing the path, wherein a context of the execution is written in a database between execution of at least some of the set of steps.

* * * * *