

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250260782

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

AGARWAL; Tanmay et al.

SPEECH-BASED VISUAL INDICATOR DURING COMMUNICATION SESSION

Abstract

A device includes one or more processors configured to detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. The one or more processors are further configured to detect that the video component includes an object that is associated with the particular speech. The one or more processors are further configured to update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

Inventors: AGARWAL; Tanmay (Najibabad, IN), SETHIA; Sandeep (Hyderabad, IN),
PONAGANTI; Ravichandra (Hyderabad, IN)

Applicant: QUALCOMM Incorporated (San Diego, CA)

Family ID: 90482403

Appl. No.: 19/196357

Filed: May 01, 2025

Related U.S. Application Data

parent US continuation 18172853 20230222 PENDING child US 19196357

Publication Classification

Int. Cl.: H04N5/262 (20060101); G10L15/08 (20060101); G10L15/26 (20060101)

U.S. Cl.:

CPC H04N5/262 (20130101); G10L15/08 (20130101); G10L15/26 (20130101);

Background/Summary

I. CROSS-REFERENCE TO RELATED APPLICATIONS [0001] The present application claims priority from and is a continuation of pending U.S. patent application Ser. No. 18/172,853, filed Feb. 22, 2023 and entitled “SPEECH-BASED VISUAL INDICATOR DURING COMMUNICATION SESSION,” the content of which is incorporated herein by reference in its entirety.

II. FIELD

[0002] The present disclosure is generally related to visual indicators used during a communication session.

III. DESCRIPTION OF RELATED ART

[0003] Advances in technology have resulted in smaller and more powerful computing devices. For example, there currently exist a variety of portable personal computing devices, including wireless telephones such as mobile and smart phones, tablets and laptop computers that are small, lightweight, and easily carried by users. These devices can communicate voice and data packets over wireless networks. Further, many such devices incorporate additional functionality such as a digital still camera, a digital video camera, a digital recorder, and an audio file player. Also, such devices can process executable instructions, including software applications, such as a web browser application, that can be used to access the Internet. As such, these devices can include significant computing capabilities.

[0004] Such computing devices can be used to facilitate voice and/or video communication sessions (such as conference calls or videoconferences). Computing devices that support video communication sessions often enable one or more participants to share visual content, such as via screen sharing that allows other participants to view a shared document. For example, a presenter can share a screen on which a series of slides are displayed by the presenter so that other participants of the video communication session can follow along as the presenter describes the content of the individual slides. Devices that support such communication sessions often provide an option for a participant to display a visual indicator such as a mouse pointer on the screen, enabling the participant to direct attention to one or more elements of the shared content being displayed. For example, when a presenter is describing a slide that includes a bulleted list of text items, the presenter can control the mouse pointer to hover on a particular bullet point that the presenter is discussing.

[0005] When used appropriately, such visual indicators are useful for directing the attention of participants and adding clarity to such presentations. However, a presenter may forget to update the position of the indicator when the presenter changes the discussion topic, which can cause confusion to the other participants. Additionally, having to repeatedly reposition the indicator as the presentation progresses can be distracting or burdensome to the presenter and restricts the presenter's ability to perform other actions with the presenter's hands, such as gesturing, holding notecards or other reference materials, or taking notes.

IV. SUMMARY

[0006] According to a particular aspect, a device includes one or more processors configured to detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. The one or more processors are further configured to detect that the video component includes an object that is associated with the particular speech. The one or more processors are further configured to update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0007] According to a particular aspect, a method includes detecting, at a device and during a

communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. The method also includes detecting, at the device, that the video component includes an object that is associated with the particular speech. The method further includes updating, at the device, the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0008] According to a particular aspect, a non-transitory computer-readable medium stores instructions that are executable by one or more processors to cause the one or more processors to detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. The instructions are further executable to detect that the video component includes an object that is associated with the particular speech. The instructions are further executable to update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0009] According to a particular aspect, an apparatus includes means for detecting, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. The apparatus also includes means for detecting that the video component includes an object that is associated with the particular speech. The apparatus further includes means for updating the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0010] Other aspects, advantages, and features of the present disclosure will become apparent after review of the entire application, including the following sections: Brief Description of the Drawings, Detailed Description, and the Claims.

Description

V. BRIEF DESCRIPTION OF THE DRAWINGS

[0011] FIG. 1 is a block diagram of a particular illustrative aspect of a system operable to apply a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0012] FIG. 2 is a diagram illustrating aspects associated with a communication session manager of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0013] FIG. 3 is a diagram illustrating aspects associated with applying a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0014] FIG. 4 is a diagram illustrating aspects associated with applying a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0015] FIG. 5 is a diagram illustrating aspects associated with applying a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0016] FIG. 6 is a diagram illustrating operations that may be performed by the communication session manager of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0017] FIG. 7 is a diagram illustrating operations that may be performed by the communication session manager of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0018] FIG. 8 is a diagram illustrating operations that may be performed by the communication

session manager of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0019] FIG. 9 is a diagram illustrating operations that may be performed by the communication session manager of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0020] FIG. 10 is a diagram illustrating operations that may be performed by the communication session manager of the system of FIG. 1, in accordance with some examples of the present disclosure.

[0021] FIG. 11 illustrates an example of an integrated circuit operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0022] FIG. 12 is a diagram of a mobile device operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0023] FIG. 13 is a diagram of a headset operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0024] FIG. 14 is a diagram of a wearable electronic device operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0025] FIG. 15 is a diagram of a voice-controlled speaker system operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0026] FIG. 16 is a diagram of a camera operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0027] FIG. 17 is a diagram of an extended reality headset operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0028] FIG. 18 is a diagram of a first example of a vehicle operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0029] FIG. 19 is a diagram of a second example of a vehicle operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

[0030] FIG. 20 is a diagram of a particular implementation of a method of applying a speech-based visual indicator to a video component of a communication session that may be performed by the device of FIG. 1, in accordance with some examples of the present disclosure.

[0031] FIG. 21 is a block diagram of a particular illustrative example of a device that is operable to facilitate application of a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure.

VI. DETAILED DESCRIPTION

[0032] Although the use of visual indicators such as mouse pointers is helpful for directing the attention of participants and adding clarity to presentations, such as during a video conference, such visual indicators can introduce confusion when a presenter forgets to update the position of the indicator. Also, having to repeatedly reposition the indicator to follow the presentation discussion can be distracting or burdensome to the presenter, in addition to restricting the presenter's ability to perform other actions with the presenter's hands.

[0033] Devices and methods are disclosed for applying a speech-based visual indicator to a video component of a communication session. The disclosed techniques provide a presenter with an

automated way to hover a pointer or otherwise highlight content displayed during a screen sharing session. According to an aspect, the presenter's speech is analyzed to identify a particular topic being discussed, and the visual indicator is automatically positioned to highlight or emphasize text or other objects in the displayed content that correspond to the identified discussion topic. By automatically updating the position of the indicator based on the content of the presenter's speech, the above-described drawbacks associated with conventional presentation systems are avoided.

[0034] According to some aspects, machine learning techniques, such as neural networks, are employed to determine the discussion topic of the presenter's speech as well as to analyze the accompanying visual content, such as a shared document, to determine where to position the visual indicator. In a particular implementation, speech topic detection is performed using speech-to-text analysis, and the visual content is analyzed using frame-to-text analysis. As the discussion topic changes, or as the visual content is updated, the location of the visual indicator is updated to remain relevant to the most recently identified discussion topic.

[0035] According to some aspects, when the discussion topic identified from the presenter's speech is not identified in the currently displayed visual content, a search can be automatically performed for relevant images associated with the discussion topic, and an image obtained via the search can be automatically inserted into the displayed content. According to another aspect, when the discussion topic identified from the presenter's speech is not identified in the currently displayed visual content, such as when the presenter begins discussing a block diagram or flowchart that is not present in the currently displayed visual content, the block diagram or flowchart can be automatically generated on-the-fly and inserted into the displayed visual content based on the presenter's description.

[0036] According to some aspects, the presenter's speech is further analyzed to detect jargon, such as one or more utterances that are associated with particular actions or usages, such as "go to next slide" or "on the left side of the screen." When such jargon is detected, the actions or usages associated with the jargon can be automatically performed or applied. In some implementations, when performing an action such as navigating to a next slide of a slide deck that is being shared, a current position of the visual indicator is recorded so that, if the presenter later returns to the current slide, the visual indicator can be automatically restored to its most recent position on the slide.

[0037] According to some aspects, the disclosed techniques can be performed based on received audio content and video content of a communication session, and thus can be performed at various devices that may be participating in the communication session. For example, the automatic insertion of visual indicators based on a presenter's speech can be performed at an end-user device of the participant that is currently speaking, at an end-user device of a participant that is currently sharing the video content, at an end-user device of a participant that is receiving both the audio and the video content from remote sources, or at a communication server that manages the communication session, as illustrative, non-limiting examples.

[0038] A technical advantage of automatically generating a speech-based visual indicator is that a presenter of the communication session can engage in a hands-free presentation that automatically provides visual emphasis to portions of the displayed slide(s) or page(s) that are currently being discussed. As a result, the presenter's hands are free to engage in other activities during the presentation while still retaining the benefits associated with a manually operated pointer. Further, the automatic positioning of the visual indicator can provide a more consistently reliable guide for viewers as compared to manually operated pointers. In addition, because the disclosed techniques can be performed based on received audio content and video content of a communication session, the automatic speech-based positioning of the visual indicator can be performed at a device that does not have access to the source documents being displayed, which can improve the functioning of a presenter's device by offloading the computational load of speech analysis and video analysis to another device, such as a remote server.

[0039] Particular aspects of the present disclosure are described below with reference to the drawings. In the description, common features are designated by common reference numbers. As used herein, various terminology is used for the purpose of describing particular implementations only and is not intended to be limiting of implementations. For example, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. Further, some features described herein are singular in some implementations and plural in other implementations. To illustrate, FIG. 1 depicts a device **102A** including one or more processors (“processor(s)” **190** of FIG. 1), which indicates that in some implementations the device **102A** includes a single processor **190** and in other implementations the device **102A** includes multiple processors **190**. For ease of reference herein, such features are generally introduced as “one or more” features and are subsequently referred to in the singular or optional plural (as indicated by “(s)”) unless aspects related to multiple of the features are being described.

[0040] In some drawings, multiple instances of a particular type of feature are used. Although these features are physically and/or logically distinct, the same reference number is used for each, and the different instances are distinguished by addition of a letter to the reference number. When the features as a group or a type are referred to herein e.g., when no particular one of the features is being referenced, the reference number is used without a distinguishing letter. However, when one particular feature of multiple features of the same type is referred to herein, the reference number is used with the distinguishing letter. For example, referring to FIG. 1, multiple devices are illustrated and associated with reference numbers **102A**, **102B**, and **102C**. When referring to a particular one of these devices, such as a device **102A**, the distinguishing letter “A” is used. However, when referring to any arbitrary one of these devices or to these devices as a group, the reference number **102** is used without a distinguishing letter.

[0041] As used herein, the terms “comprise,” “comprises,” and “comprising” may be used interchangeably with “include,” “includes,” or “including.” Additionally, the term “wherein” may be used interchangeably with “where.” As used herein, “exemplary” indicates an example, an implementation, and/or an aspect, and should not be construed as limiting or as indicating a preference or a preferred implementation. As used herein, an ordinal term (e.g., “first,” “second,” “third,” etc.) used to modify an element, such as a structure, a component, an operation, etc., does not by itself indicate any priority or order of the element with respect to another element, but rather merely distinguishes the element from another element having a same name (but for use of the ordinal term). As used herein, the term “set” refers to one or more of a particular element, and the term “plurality” refers to multiple (e.g., two or more) of a particular element.

[0042] As used herein, “coupled” may include “communicatively coupled,” “electrically coupled,” or “physically coupled,” and may also (or alternatively) include any combinations thereof. Two devices (or components) may be coupled (e.g., communicatively coupled, electrically coupled, or physically coupled) directly or indirectly via one or more other devices, components, wires, buses, networks (e.g., a wired network, a wireless network, or a combination thereof), etc. Two devices (or components) that are electrically coupled may be included in the same device or in different devices and may be connected via electronics, one or more connectors, or inductive coupling, as illustrative, non-limiting examples. In some implementations, two devices (or components) that are communicatively coupled, such as in electrical communication, may send and receive signals (e.g., digital signals or analog signals) directly or indirectly, via one or more wires, buses, networks, etc. As used herein, “directly coupled” may include two devices that are coupled (e.g., communicatively coupled, electrically coupled, or physically coupled) without intervening components.

[0043] In the present disclosure, terms such as “determining,” “calculating,” “estimating,” “shifting,” “adjusting,” etc. may be used to describe how one or more operations are performed. It should be noted that such terms are not to be construed as limiting and other techniques may be utilized to perform similar operations. Additionally, as referred to herein, “generating,”

“calculating,” “estimating,” “using,” “selecting,” “accessing,” and “determining” may be used interchangeably. For example, “generating,” “calculating,” “estimating,” or “determining” a parameter (or a signal) may refer to actively generating, estimating, calculating, or determining the parameter (or the signal) or may refer to using, selecting, or accessing the parameter (or signal) that is already generated, such as by another component or device.

[0044] FIG. 1 is a block diagram of a particular illustrative aspect of a system **100** operable to apply a speech-based visual indicator to a video component of a communication session, in accordance with some examples of the present disclosure. In FIG. 1, the system **100** includes multiple devices **102** (including device **102A**, **102B**, and **102C**), which are participating in a communication session **130**. Although FIG. 1 illustrates three devices **102**, in other implementations, the system **100** includes more or fewer devices **102**.

[0045] The communication session **130** includes at least an audio component **132** and a video component **134** that are transmitted to one or more of the devices **102**, from one or more of the devices **102**, or both, via a media stream **160** (e.g., including audio **162** and video **164**). In some implementations, the communication session **130** corresponds to at least one of a conference call, a seminar, or a multi-participant extended reality session, as illustrative, non-limiting examples. As a particular example, the communication session **130** can include a conference call that includes both audio and video components, such as a video conference with screen sharing. As explained further below, the system **100** includes one or more communication session managers **140** that are configured to automatically add a visual indicator **152**, such as a pointer indicator **154**, to the video component **134** to generate an updated video component **136** based on the speech of one or more participants of the communication session **130**. For example, when a person **110A** giving a presentation using screen sharing at the device **102A** speaks about an item on a presented slide, a communication session manager **140** can automatically add a pointer indicator **154** that points to the referenced item and that is viewable by all participants of the communication session **130**, such as a person **110B** operating the device **102B**.

[0046] In the system **100**, the devices **102** communicate via one or more networks **184**. In the example illustrated in FIG. 1, one or more communication session servers **106** of a communication service are coupled to the network **184** and operable to support the communication session **130** between the devices **102**.

[0047] FIG. 1 illustrates a particular example of aspects of the device **102A**. While details of the other devices **102B**, **102C** are not shown in FIG. 1, each of the other devices **102B**, **102C** may include similar or identical features to those described with reference to the device **102A**. In some examples, one or more of the other devices **102B**, **102C** may include additional features, fewer features, one or more different features, or a combination thereof, than those described with reference to the device **102A**. In FIG. 1, the device **102A** includes communication circuitry **120**, one or more speakers **118**, and one or more microphones **116** coupled to one or more processors **190**. In addition, the device **102A** is coupled to or includes a display device **104** configured to display the updated video component **136** that includes the visual indicator **152**.

[0048] The communication circuitry **120** is configured to enable communication via the network **184**. As illustrated, the communication circuitry **120** includes a modem **126** configured to receive at least one of the audio component **132**, the video component **134**, or the updated video component **136**, from a remote device **102** or from the communication session server(s) **106**, send at least one of the audio component **132**, the video component **134**, or the updated video component **136** to a remote device **102** or to the communication session server(s) **106**, or both. The communication circuitry **120** may also include one or more other components, such as a transceiver. In a particular aspect, the communication circuitry **120** is configured to support one or more wireless communications protocols, such as a Bluetooth® communication protocol, a Bluetooth® Low-energy (BLE) communication protocol, a Zigbee® communication protocol, a Wi-Fi® communication protocol, one or more other wireless local area network protocols, or any

combination thereof (Bluetooth® is a registered trademark of Bluetooth SIG, Inc.; Zigbee® is a registered trademark of Connectivity Standards Alliance; Wi-Fi® is a registered trademark of Wi-Fi Alliance). Additionally, or alternatively, in some implementations, the communication circuitry **120** is configured to support wide-area wireless communication protocols, such as one or more cellular voice and data network protocols from a 3rd Generation Partnership Project (3GPP) standards organization. Further, in some implementations, the communication circuitry **120** is configured to support one or more wired communications protocols. For example, in such implementations, the communication circuitry **120** also includes one or more data ports, such as Ethernet ports, universal serial bus (USB) ports, etc.

[0049] The microphone(s) **116** are configured to capture particular speech **112** of the person **110A**, and the speaker(s) **118** are configured to play out sound based on the audio component **132** of the communication session **130**. Although the microphone(s) **116** and the speaker(s) **118** are illustrated in FIG. 1 as integrated within the device **102A**, in some implementations, one or more of the microphone(s) **116**, the speaker(s) **118**, or both, are external to the device **102A** and coupled to the processor(s) **190** via one or more audio ports, data ports, or other interface circuitry.

[0050] The processor(s) **190** include a communication session manager **140A** that is operable to initiate, control, support, or otherwise perform operations associated with the communication session **130**. For example, the communication session manager **140A** may include, correspond to, or be included within an end-user application associated with the communication service. In other examples, the communication session manager **140A** is a separate application that facilitates control of the device **102A** during the communication session **130** and possibly at other times. To illustrate, the communication session manager **140A** may include a media application or plug-in that interacts with the communication session server(s) **106**.

[0051] In the example illustrated in FIG. 1, particular aspects of the communication session manager **140A** are shown, including an audio analyzer **142**, a video analyzer **146**, and a video updater **150**. In some implementations, the communication session manager **140A** includes more, fewer, or different components. For example, in some implementations, the communication session manager **140A** includes a screen sharing interface, a chat interface, or other components associated with the communication service.

[0052] The audio analyzer **142** is configured to detect, during the communication session **130**, that the audio component **132** includes particular speech, such as a topic **144**, of a participant of the communication session **130**. To illustrate, in an implementation in which the person **110A** is presenting during a video conference, the speech **112** of the person **110A** is captured by the microphones **116** to generate input audio data **122**. The input audio data **122** is provided to the processor(s) **190** and used as (or included in) the audio component **132**. In such implementations, the audio analyzer **142** processes the audio component **132** to identify particular speech, such as the topic **144**. According to some aspects, the audio analyzer **142** is configured to use a speech-to-text network to process the audio component **132** to detect the particular speech, such as described further with reference to FIG. 2 and FIG. 4.

[0053] The video analyzer **146** is configured to process the video component **134** and determine one or more objects **148** that are present in the video component **134**. For example, the video component **134** can include a slide or page of a document being presented by the person **110A**, and the video analyzer **146** identifies the object(s) **148** in the slide or page. (Although “slide” and “page” are used herein for purpose of explanation, it should be understood the present disclosure is not limited to any particular document type, and thus terms such as “slide,” “page,” “sheet,” “worksheet,” “spreadsheet,” etc. that are associated with various productivity office applications and that represent portions or subdivisions of an underlying document are to be considered interchangeable and not restrictive to any particular document type.) Examples of objects **148** that can be identified include at least one of text, a bullet point, a table element, or a graphical element that includes text, as illustrative, non-limited examples.

[0054] According to some aspects, the video analyzer **146** is configured to detect that the video component **134** includes a particular object **148** that is associated with the particular speech (e.g., the topic **144** detected by the audio analyzer **142**). In an example, the video analyzer **146** determines that an object **148A** includes or is otherwise associated with text **158A**, an object **148B** includes or is otherwise associated with text **158B**, and an object **148C** includes or is otherwise associated with text **158C**. In some implementations, the video analyzer **146** is configured to use an image-to-object network to process the video component **134** to detect the object(s) **148**, the text **158** associated with the objects **148**, or both, such as described further with reference to FIG. 2 and FIG. 5.

[0055] According to some aspects, the communication session manager **140A** performs a comparison of the topic **144** to the objects **148** to determine if an object **148** matches the topic **144**. To illustrate, the communication session manager **140A** can compare the topic **144** to the text **158** of the objects **148** and determine that the text **158C** of the object **148C** matches the topic **144**. In some implementations, such comparisons are performed by the video analyzer **146**. In other implementations, the comparisons can be performed by one or more other components of the communication session manager **140A**, such as a speech-to-object matcher as described further with reference to FIG. 2.

[0056] The video updater **150** is configured to update the video component **134** to apply a visual indicator **152** to the object **148** that is determined to match the particular speech. For example, the video updater **150** updates the video component **134** (e.g., the image of the presented slide or page of a shared document) to add the pointer indicator **154** to point to the object **148C** in response to determining that the text **158C** matches the topic **144**. Although in FIG. 1 the visual indicator **152** is displayed as a pointer indicator **154** (e.g., a dot, circle, arrow, etc.), the visual indicator **152** can generally include any indicator configured to emphasize or draw a viewer's attention to a selected object **148** and can include at least one of a pointer indicator, a text effect, or highlighting, as described further with reference to FIG. 3.

[0057] The resulting updated video component **136** can be provided to the display device **104** to display to the person **110A**. In addition, the updated video component **136** can be provided to the device **102B** and the device **102C** via transmission of a media stream **160A** from the device **102A** via the network **184**. For example, the media stream **160A** can include audio **162A** corresponding to the audio component **132** (e.g., including the speech **112**) and video **164A** corresponding to the updated video component **136** (e.g., the displayed slide or page with the pointer indicator **154** inserted). In some implementations, the media stream **160A** is transmitted to the communication session server(s) **106** and distributed from the communication session server(s) **106** to the other devices **102** participating in the communication session **130**, such as the device **102B** and the device **102C**. In other examples, the devices **102** communicate according to a peer-to-peer implementation that excludes or bypasses the communication session server(s) **106**, and the media stream **160A** is sent from the device **102A** to the other devices **102** via the network **184**.

[0058] As illustrated, instances of the communication session manager **140** can optionally be implemented at one or more of the other devices **102**, such as a communication session manager **140B** at the device **102B** and a communication session manager **140C** at the device **102C**. Optionally, a communication session manager **140D** can be implemented at the communication session server(s) **106**. Technical advantages of implementing the communication session manager **140D** at the communication session server(s) **106** can include providing a similar experience at each of the devices **102**, offloading resource intensive computations to the communication session server(s) **106**, etc. Technical advantages of implementing the communication session manager **140** at a device **102** can include compatibility with existing communication session server(s) **106** or advantages of peer-to-peer communication without a central communication session server, etc. Various use cases are described below.

[0059] In a first use case, the communication session manager **140** is included in an end-user

device configured to send a media stream **160** that includes the audio component **132** and the updated video component **136** to remote device(s) **102**. To illustrate, the device **102A** can be an end-user device of the person **110A**. The person **110A** is an “audio presenter” of the communication session **130** in that the person **110A** is designated, at least temporarily, as the participant whose speech is analyzed to determine matches to objects **148** in the currently displayed slide(s) or page(s). In this first use case, the device **102A** of the audio presenter (the person **110A**) processes the speech of the audio presenter (e.g., the speech **112** of the person **110A**), determines a match to one of the objects **148**, updates the video component **134** by adding a visual indicator **152**, and sends the resulting audio component **132** and updated video component **136**, via the media stream **160A**, to the remote devices **102B** and **102C**. In this example, while the person **110A** is the audio presenter, any of the communication session manager **140B**, the communication session manager **140C**, and the communication session manager **140D** that may be included in the system **100** refrains from performing the audio and video analysis described for the communication session manager **140A** and may instead passively receive and present the audio component **132** and the updated video component **136** from the device **102A**. In some implementations, one or more of the device **102B**, the device **102C**, or the communication session server(s) **106** do not include a communication session manager **140** configured to perform the audio and video analysis described for the communication session manager **140A** and may passively receive and present the audio component **132** and the updated video component **136** from the device **102A**.

[0060] In a second use case, the communication session manager **140** is included in an end-user device **102** configured to receive, from a remote device, a media stream **160** that includes the audio component **132** and the video component **134**. In contrast to the first use case, in the second use case the analysis of the audio presenter's speech, comparison to the objects **148**, and insertion of the visual indicator **152** is not performed at the device **102** of the audio presenter and is instead performed locally at the end-user device **102** of one or more recipients of the media stream **160**. In an example in which the person **110B** is the audio presenter, the device **102B** of the person **110B** captures the speech of the person **110B** and transmits the captured speech as audio **162B** of a media stream **160B**. The media stream **160B** also includes video **164B** representing the video component **134**, such as the currently displayed slide(s) or page(s) shared by the person **110B**.

[0061] Continuing the second use case example, the media stream **160B** is received by the device **102A**, which provides the audio component **132** as output audio data **124** for playout to the person **110A** via the speakers **118**. In addition, the communication session manager **140A** processes the audio component **132** and the video component **134** as described above and, if a match is detected between the topic **144** and an object **148**, generates the updated video component **136** and sends the updated video component **136** to the display device **104** for display to the person **110A**. If no match is detected, the device **102A** outputs the video component **134** (without adding a visual indicator **152**) to the to the display device **104** for display to the person **110A**. In some implementations, the device **102A** may send data conveying the location of the visual indicator **152** to one or more of the device **102B**, the device **102C**, or the communication session server(s) **106**, which may enable the visual indicator **152** to also be displayed at one or more of the other devices **102B**, **102C**.

[0062] In a third use case, the communication session manager **140** is included in a communication session server **106** configured to send the updated video component **136** to one or more end-user devices **102** of participants of the communication session **130**. In this use case, the analysis of the audio presenter's speech, comparison to the objects **148**, and insertion of the visual indicator **152** is not performed at the device **102** of the audio presenter and is instead performed at the communication session server **106** prior to transmission to one or more participants of the communication session **130**. In an example of the third use case in which the person **110B** is the audio presenter, the device **102B** captures the speech of the person **110B** and transmits the captured speech (e.g., the audio component **132**) to the communication session server **106** as audio of a media stream. The media stream also includes video representing the video component **134**, such as

the currently displayed slide(s) or page(s) shared by the person **110B** or shared by one or more other participants of the communication session **130**, such as the person **110A**.

[0063] Continuing the third use case example, the communication session manager **140D** at the communication session server **106** processes the audio component **132** and the video component **134** and, if a match is detected between the topic **144** and an object **148**, generates the updated video component **136**. The communication session server **106** sends the audio component **132** and the updated video component **136** (or the video component **134** if no match is detected) to the other devices **102A** and **102C** participating in the communication session **130**, such as via a media stream **160**. In some implementations, the communication session server **106** also sends the updated video component **136** to the device **102B** so that the person **110B** can also view the visual indicator **152** that has been added by the communication session server **106** to the currently displayed slide(s) or page(s).

[0064] Although depicted as laptop or desktop-type computing devices, it should be understood that one or more of the devices **102** may be another type of device. In an illustrative example, the communication session manager **140** is integrated in a headset that includes the microphone(s) **116** coupled to the processor(s) **190**, and the microphone(s) **116** are configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators **152** during the communication session **130** while the participant is engaged in physical activity. In one particular example, the headset corresponds to an audio headset, such as described further with reference to FIG. **13**. In another particular example, the headset corresponds to an extended reality headset that includes a head-mounted display device coupled to the one or more processors **190** to enable display of the updated video component **136** to a user, such as described further with reference to FIG. **17**.

[0065] In another illustrative example, the communication session manager **140** is integrated in a vehicle that includes the microphone(s) **116** coupled to the processor(s) **190**, and the microphone(s) **116** are configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators **152** during the communication session **130** while the participant is a user or occupant of the vehicle, such as described further with reference to FIG. **18** and FIG. **19**. In other implementations, the communication session manager **140** can be integrated into other devices such as a mobile phone or tablet device (e.g., FIG. **12**), a wearable device (e.g., FIG. **14**), a wireless speaker and voice activated device (e.g., FIG. **15**), or a camera device (e.g., FIG. **16**), as illustrative, non-limiting examples.

[0066] One technical advantage of processing the audio component **132** and the video component **134** and adding the visual indicator **152** when a match is detected is that presenter of the communication session **130** can engage in a hands-free presentation that automatically provides visual emphasis to portions of the displayed slide(s) or page(s) that are currently being discussed. As a result, the presenter's hands are free to engage in other activities, such as holding notecards, gesturing, interacting with a keyboard or pointing device to search notes or perform internet searches, etc., during the presentation, while still providing the other participants with the functionality that is conventionally provided by a manually operated pointer to direct the viewers' attention. Further, the automatic insertion of the visual indicator **152** based on the topic **144** of discussion can provide a more accurate and reliable guide for viewers as compared to conventional manually operated pointers that the presenter may forget to update or which may drift from their intended locations, such as due to physical movement at the presenter's device. In addition, by processing the video component **134** to detect the objects **148**, such as when the video component **134** is received as streaming video from a remote presenter, the video can be automatically enhanced at the receiving device by inserting the visual indicator **152** without requiring access to the source documents being displayed.

[0067] FIG. **2** is a diagram of an illustrative example **200** of components and operations associated with applying a speech-based visual indicator to a video component of a communication session, in

accordance with some examples of the present disclosure. In the example **200** illustrated in FIG. 2, a communication session manager **140** (e.g., the communication session manager **140A**) includes the audio analyzer **142**, the video analyzer **146**, and the video updater **150**, and may optionally include one or more other components as described in more detail below.

[0068] The audio analyzer **142** is configured to use a speech-to-text network **202** to process audio data **206** to detect particular speech in the audio data **206**. According to an aspect, the audio data **206** includes a representation of the audio component **132**, such as the speech **112** of the person **110A**, and the speech-to-text network **202** is configured to detect particular speech in the audio component **132**. A particular example of the speech-to-text network **202** is described in further detail with reference to FIG. 4.

[0069] Optionally, in some implementations, the speech-to-text network **202** is configured to detect one or more keywords **210** or one or more topics **144**. For example the keywords **210** can correspond to a sequence of the individual words detected in the audio data **206**, while the topic **144** can correspond to an overall topic of discussion that is generated based on analysis of the keywords **210**. To illustrate, for a particular segment of speech, the keywords **210** may include the sequence of words “the next row of the table is the manager the manager is responsible for roles being executed effectively and successfully,” while the topic **144** may be “manager.”

[0070] The video analyzer **146** is configured to use an image-to-object network **204** to process video data **208** to detect one or more objects **148** in the video data **208**. According to an aspect, the video data **208** includes a representation of the video component **134**, and the image-to-object network **204** is configured to detect the one or more objects **148** in the video component **134**. Optionally, for each of the detected object(s) **148**, the image-to-object network **204** is configured to detect object content **214**, an object location **216**, an object type **218**, or a combination thereof, for that object **148**. In a particular implementation, the object content **214** includes text that is detected in the video data **208** and associated with the object **148**, such as the text **158** of FIG. 1, the object location **216** indicates a position of the object **148** (e.g., coordinate locations identifying a bounding box or other segment of the video data **208** that includes the object **148**), and the object type **218** identifies what kind of object has been detected, such as text, a bullet point, an element of a table (e.g., a table cell), or an element of a flow chart, a block diagram, a timeline, or one or more other graphical objects that can include, or be associated with, text in the video data **208**. A particular example of the image-to-object network **204** is described in further detail with reference to FIG. 5.

[0071] Optionally, in some implementations, the communication session manager **140** includes a document analyzer **222** configured to, prior to a start of the communication session **130**, analyze one or more presentation documents **220** associated with the communication session **130** to detect one or more presentation objects **224** in the presentation document(s) **220**. For example, the presentation documents **220** may include one or more slide decks, formatted text documents, spreadsheets, etc., stored locally at the audio presenter's device **102** or accessed from network storage. The document analyzer **222** can be configured to parse metadata and content of the presentation document(s) **220** to generate a record of one or more presentation object(s) **224** in advance of presenting the presentation document(s) **220** during the communication session **130**. Processing of the presentation document(s) **220** prior to the communication session **130** enables generation of data (e.g., content, location, and type) of each of the presentation object(s) **224** that coincides with the data for the objects **148** that would be extracted from the video data **208** if the presentation document(s) **220** were presented during the communication session **130**. The presentation object(s) **224** are usable in place of the objects **148** identified by the video analyzer **146** so that, during the communication session **130**, detection of an object that matches particular speech in the audio data **206** is at least partially based on a determination that at least one of the presentation object(s) **224** is associated with the particular speech.

[0072] In a particular example, the communication session manager **140** is configured to determine that the video data **208** corresponds to one or more slides or pages of the presentation document(s)

220, such as by detecting that the user (e.g., a person **110**) is sharing one or more of the presentation document(s) **220** during the communication session **130**, and in response the communication session manager **140** uses the previously determined presentation object(s) **224** and deactivates, powers down, or sets to a reduced power mode the image-to-object network **204**, the video analyzer **146**, or both. The document analyzer **222** therefore provides the technical advantage of reduced computational load and reduced power consumption as compared to operating the image-to-object network **204** of the video analyzer **146** to detect the objects **148** on-the-fly during the communication session **130**.

[0073] The communication session manager **140** includes a speech-to-object matcher **240** that is configured to perform one or more comparisons of the particular speech detected in the audio data **206** (e.g., the topic(s) **144** or the keyword(s) **210**) to the objects detected in the visual presentation accompanying the audio data **206** (e.g., the object(s) **148** or the presentation object(s) **224**). In a particular example in which the audio analyzer **142** outputs the topic(s) **144** and the video analyzer **146** outputs the object(s) **148**, the speech-to-object matcher **240** receives an indication of a most recently detected topic **144** from the audio analyzer **142** and compares the topic **144** to the content **214** of the object(s) **148**. In another particular example in which the audio analyzer **142** outputs the topic(s) **144** and the presentation object(s) **224** are available, based on determining that the user is sharing one or more slides or pages of the presentation document(s) **220** during the communication session **130**, the speech-to-object matcher **240**, receives an indication of a most recently detected topic **144** from the audio analyzer **142** and compares the topic **144** to the content of the presentation object(s) **224** associated with the one or more slides or pages currently being shared. The speech-to-object matcher **240** generates an output including the object content **214**, the object location **216**, the object type **218**, or a combination thereof, of the selected object **242** (e.g., the object **148** or the presentation object **224** that matches the particular speech) to the video updater **150**.

[0074] The video updater **150** is configured to modify the video component **134** of the communication session **130** to generate the updated video component **136** by inserting the visual indicator **152** to emphasize the selected object **242**. The updated video component **136** is output by the communication session manager **140** as updated video data **272**. In some implementations, the video updater **150** is configured to use a single type of visual indicator **152**, such as the pointer indicator **154**, and a single position relative to the location of the selected object **242**, such as positioned to the side of, above, or below the selected object **242**. In other implementations, the video updater **150** determines a type of visual indicator **152** and a location of the visual indicator **152** at least partially based on the object type **218** and the object location **216** of the selected object **242**.

[0075] As an illustrative, non-limiting example, when the object type **218** corresponds to text or a bullet point, the video updater **150** adds the visual indicator **152** as a pointer indicator **154** located to the left of the selected object **242** or centered under the selected object **242**. In another illustrative, non-limiting example, when the object type **218** corresponds to table element, the video updater **150** adds the visual indicator **152** as a text effect that modifies an appearance of the text in the table element. In another illustrative, non-limiting example, when the object type **218** corresponds to text-containing graphical element, such as a block of a block diagram or flow chart, the video updater **150** adds the visual indicator **152** as highlighting within the graphical element. Determination of how and where to apply the visual indicator **152** can be based on one or more default settings, based on one or more user-selected settings, or a combination thereof.

[0076] Optionally, in some implementations, the communication session manager **140** is configured to detect instances of jargon **212** in the audio data **206**. (As used herein, “jargon” refers to particular utterances that are associated with particular actions or usages, such as “next slide” or “on the left side.”) In such implementations, the speech-to-text network **202** is configured to detect one or more utterances **232** corresponding to jargon **212**, which are provided to an optional jargon mapper **230** for processing. To illustrate, the speech-to-text network **202** may be configured to

detect the jargon **212** by detecting, during the communication session **130**, that the audio component **132** includes an utterance **232** that is mapped to a particular action **234** associated with a presentation document **220**. For example, the jargon mapper **230** can include a list of utterances **232** that are mapped to respective actions **234**, such as a first utterance **232A** mapped to a first action **234A** and a second utterance **232B** mapped to a second action **234B**. The speech-to-text network **202** may compare the generated keywords **210** detected in the audio data **206** to a list of the utterances **232** and, when a match is detected, may provide an indication of the detected utterance **232** to the jargon mapper **230**. The jargon mapper **230** identifies the particular action **234** associated with the detected utterance **232**, and the communication session manager **140** causes the particular action **234** to be performed. In some examples, the video updater **150** updates the video component **134** to depict a result of performance of the particular action **234**. Examples of jargon utterances **232** and associated actions **234** are described in further detail with reference to FIG. 7.

[0077] Optionally, in some implementations, the communication session manager **140** includes an image retriever **280** that is configured to perform a search of one or more networks **290** to locate a diagram **284** (e.g., an image) associated with a discussion topic **282**. For example, the communication session manager **140** may determine a discussion topic **282** based on the topic **144**, the keywords **210**, or both. If the discussion topic **282** is not located on the slide or page currently being displayed (e.g., no matching object **148** or matching presentation object **224** is found), and if the slide or page is determined to be blank (e.g., devoid of any objects) or is determined to have a designated area (also referred to as a field of view (FOV) panel) that is blank, the communication session manager **140** can cause the image retriever **280** to initiate a search for a diagram related to the discussion topic **282**. In some implementations, the search is restricted to one or more content servers **292** that include a curated image collection **294** so that the retrieved diagram **284** is likely to be accurate and relevant to the discussion topic **282**. As an illustrative, non-limiting example, the search may be restricted to a Wikipedia search (Wikipedia® is a registered trademark of Wikimedia Foundation, Inc.). The retrieved diagram **284** is provided to the video updater **150**, and the video updater **150** updates the video component **134** (or the updated video component **136**) to include the diagram **284** in the updated video component **136**. Operation of the communication session manager **140** to identify and retrieve diagrams **284** based on discussion topics **282** for insertion into the video component **134** of the communication session **130** is described in further detail with reference to FIG. 8 and FIG. 9.

[0078] Optionally, in some implementations, the communication session manager **140** maintains a pointer position history **250**. The pointer position history **250** includes one or more tables or other data structures in which, for each page **252**, of a presentation document **220**, that was previously displayed during the communication session **130**, the pointer position history **250** tracks a position **254** of a most recent visual indicator **152** that was depicted on that page **252**. For example, when a first page **252A** was displayed previously, the last visual indicator **152** displayed on the first page **252A** was at a first position **254A**. Similarly, when a second page **252B** was previously displayed, the last visual indicator **152** displayed on the second page **252B** was at a second position **254B**. As a result, when the presenter returns to a previously presented page during the communication session **130**, the communication session manager **140** can access the pointer position history **250** and restore the visual indicator **152** to its previous state for that page. In an illustrative example, the presenter may state “go back to page **1**,” which may be identified as jargon **212** associated with a particular action **234** for re-displaying a particular page **252** of the previously displayed pages. In response, the communication session manager **140** may cause the video updater **150** to update the video component **134** to display the particular page **252** and to restore the most recent visual indicator **152** for the particular page **252**. However, the pointer position history **250** is not necessarily limited to use with navigation instructions detected via jargon **212** and may instead, or in addition, be used in conjunction with manual navigation between pages by the presenter.

[0079] Optionally, the communication session manager **140** includes an object generator **270**

configured to generate a representation of a presentation object on-the-fly when the presentation object is being described by the presenter but is not depicted on the currently presented page of a presentation document **220**. For example, the communication session manager **140** may detect, during the communication session **130**, that the audio component **132** includes a description of a component of a presentation object. To illustrate, the presenter may speak “let’s learn about a block diagram of a central processing unit that includes a control unit, an arithmetic logic unit, and main memory” while no block diagram is identified in the currently displayed page (e.g., the objects **148** and/or the presentation objects **224** do not include graphical elements with text that include “central processing unit,” “control unit,” “arithmetic logic unit,” or “main memory”).

[0080] Based on determining that the component (e.g., a block including the text “control unit”) of the presentation object (e.g., a block diagram of a central processing unit) is not present in the video component **134**, and that the page is determined to be blank or is determined to have a designated area that is blank, the object generator **270** signals the video updater **150** to update the video component **134** to generate, in the updated video component **136**, a representation of the component of the presentation object based on the description. In some implementations, the object generator **270** sends a description of the component to be added, such as the size and shape of the component, text to be added in the component, the position of the component, connections to other components that may already have been added, or a combination thereof. In this manner, the communication session manager **140** can generate a presentation object on-the-fly by iteratively adding components of a presentation object, such as elements of a flow chart, a timeline, or a block diagram, as the presenter is describing the components, as described further with reference to FIG. **10**.

[0081] FIG. **3** is a diagram **300** illustrating examples associated with applying a speech-based visual indicator to a video component of a communication session. In particular, FIG. **3** depicts a first page **302** and a second page **304** of a presentation document **220** that may be presented during the communication session **130** and also depicts an object database **306** (e.g., a table) that may be generated by the communication session manager **140** during the communication session **130**.

[0082] The first page **302** includes text **310A**, a table **320**, and a bullet point **330**. The text **310A** corresponds to a title of the first page **302**. The table **320** includes table elements **322** arranged in three rows and two columns. As illustrated, each row corresponds to an organizational role that can be assigned to a participant of a group. The top row includes a table element **322A** that includes the text “Coordinator” and a table element **322B** that includes a description of a coordinator role. The middle row includes a table element **322C** that includes the text “Organizer” and a table element **322D** that includes a description of an organizer role. The bottom row includes a table element **322E** that includes the text “Manager” and a table element **322F** that includes a description of a manager role. The bullet point **330** is located under the table **320** and includes additional information regarding group roles.

[0083] In a particular implementation, when the first page **302** is presented during the communication session **130**, the video analyzer **146** processes the video component **134** corresponding to the first page **302** and identifies the text **310A**, each of the six elements **322A-322F** of the table **320**, and the bullet point **330** as individual objects **148**. The communication session manager **140** populates a separate row of the object database **306** with information of a respective one of the identified objects **148**. As illustrated, a first row of the object database **306** is associated with the text **310A** and includes an object identifier (ID) **360** of “1,” an object type **218** of “text,” a page indicator **364** of “1,” an object location **216** that identifies the location of the text **310** (e.g., coordinates on the first page **302**), and an object content **214** of “Group Roles.” A second row through a seventh row of the object database **306** correspond to the table elements **322A-322F**, respectively, and an eighth row of the object database **306** corresponds to the bullet point **330**.

[0084] During the communication session **130** and while the first page **302** is being presented, the audio analyzer **142** processes the presenter’s speech received as the audio component **132** to detect

a topic **144**. For example in response to the audio component **132** including the speech “The first role we are going to discuss is the coordinator. A good coordinator can directly affect the success or failure of the group, so take care in appointing someone as coordinator. As we can see, the responsibilities of the coordinator include . . .” the audio analyzer **142** determines that the topic **144** is “coordinator.” The communication session manager **140** (e.g., the speech-to-object matcher **240**) compares the topic **144** “coordinator” to the object content **214** of one or more of the objects **148** in the object database **306** until either a match is determined or until all entries for objects **148** in the first page **302** have been checked and determined to not match the topic **144**. As illustrated, a match is determined for the object having object ID=2 and corresponding to the table element **322A**, which has the object content **214** of “coordinator.”

[0085] Although in this example the topic **144** and the object content **214** of the identified object have identical text (e.g., an exact match), in some implementations a match may be determined even though the topic **144** and the object content **214** are not identical. For example, the speech-to-object matcher **240** may determine a difference (or similarity) metric or a likelihood of a match, such as by direct calculation or using a network trained to estimate the difference (or similarity) metric or the likelihood, and may compare the difference (or similarity) metric or the likelihood to a threshold to determine whether a match is detected.

[0086] In response to determining the match, the video updater **150** applies the visual indicator **152** to the identified object by applying a text effect **390** to the text in the table element **322A**. As illustrated, applying the text effect **390** includes modifying the text in the table element **322A** to be boldfaced and italicized. In other examples, the text effect **390** can include other types of effects, such as changing underlining, text color, text size, font, or a combination thereof.

[0087] After discussing the first page **302**, the presenter may proceed to another page in the presentation document **220** so that the second page **304** is presented. The second the second page **304**. The diagram **340** is a block diagram that includes multiple components (blocks), each of which corresponds to a graphical element (e.g., a rectangle) that includes text. As illustrated, a graphical element **342A** includes the text “Input,” a graphical element **342B** includes the text “CPU,” a graphical element **342C** includes the text “Output.” Graphical elements **342D**, **342E**, and **342F** include the text “Control Unit,” “ALU,” and “Main Memory,” respectively.

[0088] In a particular implementation, when the second page **304** is presented during the communication session **130**, the video analyzer **146** processes the video component **134** corresponding to the second page **304** and identifies the text **310B** and each of the six graphical elements **342A-342F** of the diagram **340** as individual objects **148**. The communication session manager **140** populates a separate row of the object database **306** with information of a respective one of the identified objects **148** of the second page **304**. As illustrated, a ninth row of the object database **306** corresponds to the text **310B**, and a tenth through a fifteenth row of the object database **306** correspond to the graphical elements **342A-342F**, respectively.

[0089] During the communication session **130** and while the second page **304** is being presented, the audio analyzer **142** processes the presenter's speech received as the audio component **132** to detect a topic **144**. For example in response to the audio component **132** including the speech “The CPU includes a control unit, and the control unit coordinates operations among the other component. For example, the control unit is responsible for . . .” the audio analyzer **142** determines that the topic **144** is “control unit.” The communication session manager **140** (e.g., the speech-to-object matcher **240**) compares the topic **144** “control unit” to the object content **214** of one or more of the objects **148** in the object database **306** until either a match is determined or until all entries for objects **148** in the second page **304** have been checked and determined to not match the topic **144**. As illustrated, a match is determined for the object having object ID=13 and corresponding to the graphical element **342A**, which has the object content **214** of “control unit.”

[0090] In response to determining the match, the video updater **150** applies the visual indicator **152** to the identified object by applying highlighting **392** to the text in the graphical element **342D**. As

illustrated, applying highlighting **392** includes modifying a color of the area surrounding the text in the graphical element **342D** to be visually distinct, such as by applying yellow highlighting **392**. [0091] FIG. **4** depicts an illustrative example **400** of the speech-to-text network **202** that may be included in the communication session manager **140**. The speech-to-text network **202** includes a feature extractor **402** configured to process the audio data **206**, such as using normalization, windowing, etc., and to generate feature data associated with the audio data **206**, such as a spectrogram **404** (e.g., a log spectrogram or mel scale spectrogram). An acoustic model **406** is configured to process the feature data (e.g., the spectrogram **404**) to determine the probability distributions $P_t(c)$ over vocabulary characters 'c' per each time step 't'. As an illustrative, non-limiting example, the acoustic model **406** can be implemented using a Jasper or Quartznet-type architecture.

[0092] A decoder **410** processes the output of the acoustic model **406** (e.g., the $P_{\text{sub.t}}(c)$) to select letters, such as based on highest probability. A language model **412** is configured to guide the selection at the decoder **410** based on likelihoods of certain words appearing in context. The decoder **410** can thus generate text corresponding to the keywords **210** identified in the audio data **206** using automatic speech recognition processing techniques. A topic detector **420** is configured to process the keywords **210** to determine the topic(s) **144** associated with the audio data **206**. Optionally, in some implementations, the speech-to-text network **202** also includes a jargon detector **430** configured to process the keywords **210** to identify one or more utterances that correspond to instances of jargon **212**.

[0093] In some implementations, one or more (or all) of the feature extractor **402**, the acoustic model **406**, the decoder **410**, the language model **412**, the topic detector **420**, or the jargon detector **430** is implemented as a trained model, such as a neural network. According to an aspect, the functionality associated with two or more of the feature extractor **402**, the acoustic model **406**, the decoder **410**, the language model **412**, the topic detector **420**, or the jargon detector **430** can be combined into a single component. For example, in some implementations, the functionality described for the topic detector **420**, the jargon detector **430**, or both, can be included in the decoder **410**.

[0094] FIG. **5** depicts an illustrative example **500** of the image-to-object network **204** that may be included in the communication session manager **140**. The image-to-object network **204** includes a pre-processor **510** configured to process an input image **502** (e.g., a frame of a page included in the video data **208**) and a detection model **520** configured to process the output of the pre-processor **510**. The pre-processor **510** includes a greyscale converter **512** configured to convert the input image **502** to greyscale, an image enhancer **514** configured to perform one or more enhancement operations (e.g., contrast enhancement, noise removal) on the greyscale image to generate an enhanced greyscale image, and a binary converter **516** configured to perform a binary conversion operation on the enhanced greyscale image to generate a binary image that is output by the pre-processor **510**.

[0095] The detection model **520** can correspond to a convolutional neural network (CNN) and, according to some aspects, is based on a ResNet18 architecture. The detection model **520** is configured to process the binary image using a convolutional neural network **522** to generate feature maps **524** associated with features of the binary image. The feature maps **524** may include one or more regions of interest (RoI) corresponding to segments of the binary image that include text. RoI pooling **526** may be performed to enable classification of arbitrary-sized regions of interest at one or more classifiers **528** configured to identify text. Each segment of the image that includes text can correspond to an object **148**, where the identified text corresponds to the object content **214** and the segment position corresponds to the object location **216**, and one or more of the classifiers **528** may be configured to detect the object type **218** associated with the segment. In an illustrative example, the detection model **520** can be trained using as input a set of images containing text to be recognized, and generating the sequence of characters of the detected text as

the object content **214**. The detection model **520** can further be trained to detect the location of an image segment that includes the detected text the object location **216**, to detect the object type **218**, or both. High accuracy performance of the detection model **520** can also be achieved using long short-term memory (LSTM) units.

[0096] FIG. **6** depicts a flowchart of a particular implementation of a method **600** including operations that may be performed by a communication session manager **140**.

[0097] The method **600** includes, at block **602**, determining that a conference call session is initiated. For example, the conference call session may correspond to a particular instance of the communication session **130**.

[0098] The method **600** includes, at block **604**, analyzing audio content of a speaker (e.g., an audio presenter of the conference call session) and screen share content. For example, the audio content can correspond to the audio component **132** that is analyzed by the audio analyzer **142**, and the screen share content can correspond to the video component **134** that is analyzed by the video analyzer **146**.

[0099] The method **600** includes, at block **606**, extracting a topic from the audio content. For example, the audio analyzer **142** may process the audio component **132** to extract the topic **144**.

[0100] The method **600** includes, at block **608**, applying frame-to-text analysis and determining one or more topics associated with the frame. For example, a frame of the video data **208** can include screen share content that corresponds to a page of a presentation document **220** and that is represented in the video component **134**. The video component **134** is processed by the video analyzer **146** to determine the object(s) **148** associated with the frame. In other implementations, the one or more topics can be determined based on the presentation object(s) **224**.

[0101] The method **600** includes, at block **610**, determining whether the speaker topic matches the one or more frame topics. For example, the speech-to-object matcher **240** compares the topic **144** (from the audio content) to the object(s) **148** (or presentation object(s) **224**) to see if the text associated with any of the object(s) **148** matches the topic **144**. If no match is detected, the method **600** returns to analyzing the audio content and the screen share content at block **604**.

[0102] Otherwise, when a match is detected, the method **600** includes, at block **612**, hovering a mouse pointer on the matched topic/context. For example, the video updater **150** generates the updated video component **136** by inserting the pointer indicator **154** at or near the text of the matching object **148**.

[0103] The method **600** includes, at block **614**, determining whether the screen share and audio of the conference call session are still ongoing. In response to determining that the screen share and audio continuing are ongoing, the method **600** returns to analyzing the audio content and the screen share content at block **604**, thus continuing analysis while the speech and screen share are active. Otherwise, the screen sharing session is deactivated, at block **616**.

[0104] FIG. **7** is a diagram that includes a flowchart of a method **700** including operations that may be performed by a communication session manager **140** according to a particular implementation in which jargon detection is implemented.

[0105] The method **700** includes, at block **702**, determining that a conference call session is initiated. For example, the conference call session may correspond to a particular instance of the communication session **130**.

[0106] The method **700** includes, at block **704**, analyzing audio content of a speaker (e.g., an audio presenter of the conference call session) and screen share content. For example, the audio content can correspond to the audio component **132** that is analyzed by the audio analyzer **142**, and the screen share content can correspond to the video component **134** that is analyzed by the video analyzer **146**.

[0107] The method **700** includes, at block **706**, extracting keywords from the audio content. For example, the audio analyzer **142** may process the audio component **132** to extract the keywords **210**.

[0108] The method **700** includes, at block **708**, determining whether the keywords match any jargon. For example, the keywords **210** can be compared to a jargon database **750** to determine whether the keywords **210** match identified jargon that is associated with one or more particular actions or usages. The jargon database **750** includes multiple rows each corresponding to a respective utterance (or set of keywords) that is categorized as jargon, including values for a jargon ID **760**, jargon **762**, and an action/usage **764** of the jargon **762**. According to an aspect, the jargon database **750** corresponds to the jargon mapper **230** of FIG. 2.

[0109] In response to detecting that the keywords match the jargon with little or no ambiguity (e.g., a “straight hit”), the method **700** advances to block **714**. As a particular example in which the speaker has been discussing a particular bullet point in a displayed slide (e.g., the shared content) and a mouse pointer has been hovering at the particular bullet point, the keywords **210** can include the sequence of words “in the next point we see that . . .” These keywords, when used in a search of the jargon database **750**, result in a non-ambiguous match with the jargon “next point” corresponding to the second row of the jargon database **750** and associated with the action or usage of hovering a mouse pointer at a next sequential bullet point in a displayed slide (e.g., the screen share content).

[0110] Another possibility is that the keywords are detected to match the jargon, but with more than a threshold amount of ambiguity. For example, the keywords **210** can include the sequence of words “in the next interesting point we see that . . .” which likely matches the jargon “next point” corresponding to the second row of the jargon database **750**, but with sufficient ambiguity that the communication session manager **140** cannot predict with certainty that the next sequential bullet point in the slide is the “next interesting point” mentioned by the speaker. In such cases, the method **700** moves to block **710** for topic extraction. Otherwise, when no jargon match is detected for the keywords, jargon processing for that set of keywords ends, although non-jargon topic extraction and matching can continue in a similar manner as described in FIG. 6.

[0111] To illustrate, the method **700** includes, at block **710**, extracting a topic using the detected keywords. For example, extracting the topic can be performed by the audio analyzer **142**, such as at the topic detector **420** of the speech-to-text network **202** of FIG. 4. In the event that a jargon match with ambiguity was determined at block **708**, information regarding the potential jargon match can be used to inform the topic extraction, the topic extraction can be used to resolve the ambiguous jargon match, or both. Continuing the above example in which it is unclear whether the next sequential bullet point in the slide is the “next interesting point,” extraction of a topic that matches the text of the next sequential bullet point can confirm the jargon match. As another example, if the topic extraction is ambiguous, such as when the extracted topic can potentially match the next sequential bullet point or another object elsewhere on the slide (e.g., the text in a table or a block diagram), the potential jargon match with the next sequential bullet point can be used to resolve the ambiguity in favor of the next sequential bullet point.

[0112] The method **700** includes, at block **712**, applying frame-to-text analysis and determining one or more topics associated with the frame. For example, a frame of the video data **208** can include screen share content that corresponds to a slide of a presentation document **220** and that is represented in the video component **134**. The video component **134** is processed by the video analyzer **146** to determine the object(s) **148** associated with the frame. In other implementations, the one or more topics can be determined based on the presentation object(s) **224**.

[0113] The method **700** includes, at block **714**, determining whether relevant topic/context is on screen. For example, if the mouse pointer is currently hovering at a bullet point and the block **174** is reached via a “straight hit” jargon match to hover at a next sequential bullet point, a determination can be made as to whether a next sequential bullet point is present on screen via searching the detected frame topic(s) (e.g., the object(s) **148**) for a next sequential bullet point, such as a bullet point located immediately below the current bullet point. In such cases, the determination at block **714** may be performed without consideration of the speaker topic (if any)

that may have been extracted based on the keywords. In another example, the determination may be based on a jargon match that was resolved using the extracted speaker topic, as explained previously, and can involve comparison of the action/usage associated with the jargon with the extracted frame topic(s) (e.g., the object(s) **148** or the presentation object(s) **224**). In another example the determination may be based on the extracted speaker topic, such as when no jargon match was found, such as by comparing the speaker topic to the frame topic(s), such as via the speech-to-object matcher **240** comparing the topic **144** (from the audio content) to the object(s) **148** (or the presentation object(s) **224**) to see if the text associated with any of the object(s) **148** (or the presentation object(s) **224**) matches the topic **144**.

[0114] If the relevant topic/context is not detected on screen, the method **700** returns to analyzing the audio content and the screen share content at block **704**. Otherwise, when the relevant topic/context is detected on screen, the method **700** includes, at block **716**, hovering a mouse pointer on the matched topic/context. For example, the video updater **150** generates the updated video component **136** by inserting the pointer indicator **154** at a position at or near the text of the selected object **242** (e.g., a next selected bullet point, or other text matching the speaker topic).

[0115] The method **700** includes, at block **718**, determining whether the screen share and audio of the conference call session are still ongoing. In response to determining that the screen share and audio continuing are ongoing, the method **700** returns to analyzing the audio content and the screen share content at block **704**, thus continuing analysis while the speech and screen share are active. Otherwise, the screen sharing session is deactivated, at block **720**.

[0116] FIG. **8** is a diagram illustrating operations **800** that may be performed by a communication session manager **140** according to a particular implementation in which image retrieval is supported.

[0117] The operations **800** include detecting that a field of view panel **804** of a presented screen **802** is empty (e.g., blank, devoid of text or other objects **148**). For example, a user starts presenting content (e.g., a slide of a presentation document **220**) during a communication session **130**, and the video analyzer **146** determines that the field of view panel **804** is devoid of any text or other objects **148** during processing of the video component **134**, or the document analyzer **222** may be configured to make the determination during processing of the presentation document **220** (prior to a start of the communication session **130**). In some implementations, the field of view panel **804** corresponds to a defined region encompassing a central portion of the presented screen and excluding peripheral regions that may include header text, footer text, a title, etc., of a presented slide. Although a single field of view panel is shown, according to some aspects, the presented screen **802** includes multiple field of view panels **804** that may each be evaluated for the presence of text or other objects **148**.

[0118] The operations **800** include converting a speaker's speech to text, at block **806**. For example, the image-to-object network **204** of the audio analyzer **142** extracts the keywords **210**. As an illustrative example, the user may begin talking about the topic "fast Fourier transform," and the user's speech is converted to text (e.g., the keywords **210**).

[0119] The operations **800** include extracting a topic based on the text, at block **808**. For example, the topic detector **420** processes the keywords **210** to extract the topic **144**. Continuing the illustrative example in which the user begins talking about fast Fourier transforms, the text corresponding to the user's speech is processed to extract the topic "fast Fourier transform."

[0120] The operations **800** include, once the topic is extracted, determining whether one or more objects associated with the extracted topic is present on the presented screen **802** and, when there is not an associated object (e.g., a diagram of a fast Fourier transform) detected on the screen **802**, automatically retrieving a relevant image from a network (e.g., the internet) and inserting the retrieved image into the empty field of view panel **804**, upon which the user can continue explaining the image. To illustrate, the image retriever **280** may search the content server(s) **292** for an image related to the discussion topic **282** (e.g., "fast Fourier transform"). A retrieved image **812**

(e.g., a block diagram of an example fast Fourier transform algorithm structure) is provided to the video updater **150** for insertion into the empty field of view panel **804**, resulting in a modified screen **814**.

[0121] FIG. **9** is a diagram that includes a flowchart of a method **900** including operations that may be performed by a communication session manager **140** according to a particular implementation in which image retrieval is supported.

[0122] The method **900** includes, at block **902**, determining that a conference call session is initiated. For example, the conference call session may correspond to a particular instance of the communication session **130**.

[0123] The method **900** includes, at block **904**, analyzing audio content of a speaker (e.g., an audio presenter of the conference call session). For example, the audio content can correspond to the audio component **132** that is analyzed by the audio analyzer **142**.

[0124] The method **900** includes, at block **906**, detecting keywords in the audio content. For example, the audio analyzer **142** may process the audio component **132** to extract the keywords **210**.

[0125] The method **900** includes, at block **908**, determining whether a topic is identified based on the detected keywords. For example, the audio analyzer **142** (e.g., the topic detector **420** of the speech-to-text network **202** of FIG. **4**) may process the keywords **210** to determine whether a topic **144** is identified. In the event that a topic is not identified, the method returns to block **904**, where analysis of audio content continues to be performed.

[0126] Otherwise, when a determination is made that a topic is identified, at block **908**, the method **900** includes checking for the field of view, at block **910**. Checking for the field of view can include detecting whether a field of view panel in a displayed page is empty, such as described for the field of view panel **804** of the presented screen **802** of FIG. **8**. In an illustrative example, the video analyzer **146** determines whether the field of view panel **804** is devoid of any text or other objects **148** during processing of the video component **134**, or the document analyzer **222** may be configured to make the determination during processing of the presentation document **220**.

[0127] The method **900** includes, at block **912**, searching for relevant content. For example, the image retriever **280** may search the content server(s) **292** for content related to the identified topic (e.g., the discussion topic **282**) after determining that the displayed page does not include content matching the identified topic.

[0128] The method **900** includes, at block **914**, inserting content from the search results. For example the image retriever **280** may receive a diagram **284** as a result of the search, which is provided to the video updater **150** for insertion into the empty field of view panel.

[0129] FIG. **10** is a diagram that illustrates operations that may be performed by a communication session manager **140** according to a particular implementation in which on-the-fly generation of presentation objects is supported.

[0130] For example, during a communication session **130**, a user can present a slide and begin discussing an item that is not on the slide. For example, the user's speech can include a description **1000** of one or more components of a presentation object, such as a flow chart, a timeline, or a block diagram (as illustrative, non-limiting examples) although the described presentation object is not included in the displayed slide, and instead the slide has an empty field of view panel. The illustrated description **1000** depicts an example of words spoken by the user while the slide is displayed, and may correspond to keywords that are determined based on processing the user's speech.

[0131] The communication session manager **140** may perform a method **1020** that includes, at block **1022**, determining that a conference call session is initiated. For example, the conference call session may correspond to a particular instance of the communication session **130**.

[0132] The method **1020** includes, at block **1024**, analyzing audio content of a speaker (e.g., an audio presenter of the conference call session). For example, the audio content can include the

description **1000** and may correspond to the audio component **132** that is analyzed by the audio analyzer **142**.

[0133] The method **1020** includes, at block **1026**, detecting keywords in the audio content. For example, the audio analyzer **142** may process the audio component **132** to extract keywords **210** corresponding to (e.g., matching) the description **1000**.

[0134] The method **1020** includes, at block **1028**, determining whether a presentation object component is identified in the audio content. For example, the audio analyzer **142** or the object generator **270** may be configured to detect whether the keywords **210** include one or more terms that indicate that the speaker is describing a presentation object, such as the terms “flow chart,” “timeline,” or “block diagram,” as illustrative, non-limiting examples. As an illustrative example, the implementation **1100** includes the term “block diagram” indicating a presentation object **1010**.

[0135] When one or more terms that indicate a description of a presentation object are detected, the audio analyzer **142** or the object generator **270** process the keywords **210** to detect whether the keywords **210** include one or more terms that indicate a component of the presentation object. As an illustrative example, the component of the presentation object includes an element of a flow chart, a timeline, or a block diagram. In a particular implementation, detection of the one or more terms includes processing the keywords **210** using a trained machine learning model that is configured to identify presentation object components. For example, the implementation **1100** includes the term “central processing unit CPU” corresponding to a component **1012A**, the term “control unit” corresponding to a component **1012B**, the term “arithmetic logic unit ALU” corresponding to a component **1012C**, the term “main memory” corresponding to a component **1012D**, and the term “output” corresponding to a component **1012F**.

[0136] In some implementations, the audio analyzer **142** or the object generator **270** are also configured to detect whether the keywords **210** include one or more terms that indicate component examples, such as by processing the keywords using a trained machine learning model configured to identify examples of identified presentation object components. For example, the description **1000** includes the terms “mouse” and “keyboard” corresponding to an example **1014A** and an example **1014B**, respectively, of the “input” component **1012E**. The description **1000** also the terms “monitor” and “printer” corresponding to an example **1014C** and an example **1014D**, respectively, of the “output” component **1012F**. In some implementations, the audio analyzer **142** or the object generator **270** includes a trained machine learning model configured to process the keywords **210** and generate an output indicating each presentation object, presentation object component, and component example in the keywords **210**.

[0137] In response to determining that one or more presentation object components **1012** have been identified, at block **1028**, the method **1020** advances to block **1030**. Otherwise, the method **1020** returns to analyzing audio content of the speaker, at block **1024**.

[0138] The method **1020** includes, at block **1030**, determining whether the presentation object is on screen. For example, the video analyzer **146** may process the video component **134** to determine whether the identified presentation object (e.g., a block diagram corresponding to a CPU) is included in the slide being presented. If the presentation object is not on screen, at block **1030**, the method **1020** includes adding the presentation object on screen, at block **1032**. For example, the object generator **270** may instruct the video updater **150** to initialize generation of the identified presentation object in an empty field of view panel of the currently presented slide, such as by adding a title (e.g., “block diagram”), a border to enclose the block diagram, one or more initial blocks (e.g., one or more empty blocks), or a combination thereof.

[0139] The method **1020** includes, after determining that the presentation object is on-screen, at block **1030**, or after adding the presentation object, at block **1032**, determining whether each identified component of the presentation object is on screen, at block **1034**. For example the video analyzer **146**, the object generator **270**, or both, may be configured to determine, for each of the identified components **1012**, whether that component has already been added to the presentation

object on the screen. If all of the presentation object components identified in the keywords **210** are determined to have been added to the presentation object on the screen (e.g., the block diagram includes a block for each of the components **1012** identified in the keywords **210**), the method **1020** returns to analyzing next audio content of the speaker, at block **1024**.

[0140] Otherwise, when one or more of the presentation object components identified in the keywords **210** are determined to have not been added to the presentation object on the screen (e.g., the block diagram does not include a block for each of the components **1012** identified in the keywords), the method **1020** includes, at block **1036**, adding each such component to the screen. For example, the object generator **270** causes the video updater **150** to add each missing component **1012** to the presentation object **1010**.

[0141] FIG. **10** also depicts an illustrative example of the presentation object **1010** that may be generated on-the-fly by the communication session manager **140** based on the speaker's description **1000**. As illustrated, the presentation object **1010** is arranged as a block diagram with each of the components **1012** depicted as a block that includes text indicating the keyword(s) corresponding to that component **1012** (e.g., the component **1012A** includes the keyword “CPU” and the component **1012E** includes the keyword

[0142] “Input”). In addition, text corresponding to each of the examples **1014** is added to the corresponding component **1012** (e.g., the component **1012E** also includes the keywords “mouse” and “keyboard” corresponding to the examples **1014A** and **1014B**, respectively). Each of the blocks of the block diagram is either coupled to (e.g., via a connector line) or nested inside of another block of the block diagram, based on relationships between the components **1012** that may be determined via analysis of the description **1000**.

[0143] Thus, performance of the method **1020** enables the communication session manager **140** detect, during the communication session **130**, that the audio component **132** includes the description **1000** of a component **1012** of a presentation object **1010** and, based on determining that the component **1012** of the presentation object **1010** is not present in the video component **134**, update the video component **134** to generate a representation of the component **1012** of the presentation object **1010** based on the description **1000**.

[0144] FIG. **11** depicts an implementation **1100** in which an integrated circuit **1102** includes the one or more processors **190** of the device **102A** of FIG. **1**. The integrated circuit **1102** also includes a signal input **1104**, such as one or more bus interfaces, to receive input data **1106** for processing. For example, the input data **1106** may include data from the communication circuitry **120** or the speaker(s) **118** of FIG. **1**, such as the media stream **160** representing the audio component **132** and the video component **134**, the input audio data **122** representing the speech **112** of the person **110A**, or a combination thereof.

[0145] The integrated circuit **1102** also includes a signal output **1108**, such as a bus interface, to enable sending of output data **1110**. For example, the output data **1110** may include data provided by the processor(s) **190** to one or more of the communication circuitry **120**, the speaker(s) **118**, or the display device **104** of FIG. **1**, such as the output audio data **124**, the media stream **160** representing the audio component **132** and the video component **134** or the updated video component **136**, the visual indicator **152**, other data associated with the communication session **130**, or a combination thereof.

[0146] FIG. **12** depicts an implementation **1200** in which one of the devices **102** of FIG. **1** is a mobile device **1202**, such as a phone or tablet, as illustrative, non-limiting examples. The mobile device **1202** includes the microphone(s) **116**, the speaker(s) **118**, and a display screen **1204** (e.g., the display device **104**). A communication session manager **140** is integrated in the mobile device **1202** and illustrated using dashed lines to indicate internal components that are not generally visible to a user of the mobile device **1202**.

[0147] In a particular example, the mobile device **1202** is operable to automatically apply a speech-based visual indicator to a video component **134** of a communication session **130**. To illustrate,

when a user of the mobile device **1202** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document that the user is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at the display screen **1204**. The mobile device **1202** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0148] In implementations in which the user of the mobile device **1202** is not the audio presenter, such as when the mobile device **1202** receives the presenter's speech via a media stream, the communication session manager **140** can process the presenter's speech, identify the text or other visual element that the presenter is speaking about, and add the pointer indicator **154** or other visual indicator **152** for display at the display screen **1204** while the presenter's speech is played out via the speaker(s) **118**.

[0149] FIG. **13** depicts an implementation **1300** in which one of the devices **102** of FIG. **1** is a headset device **1302**. Components of the processor **190**, including the communication session manager **140**, are integrated in the headset device **1302**. The headset device **1302** includes the microphone **116** and the speaker(s) **118**. The microphone **116** is configured to capture the speech of a user of the headset device **1302** to enable the user to deliver a hands-free presentation with automatic visual indicators **152** during a communication session **130**, such as while the user is engaged in physical activity. Although the headset device **1302** is depicted as an over-ear type headset device, in other implementations the headset device **1302** can be configured as an in-ear type headset device, such as a pair of earbuds.

[0150] In a particular example, the headset device **1302** is operable to automatically apply a speech-based visual indicator to a video component **134** of a communication session **130**. To illustrate, when a user of the headset device **1302** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone **116** and processed to identify text (or some other element) of a shared document that the user is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. The headset device **1302** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0151] FIG. **14** depicts an implementation **1400** in which one of the devices **102** of FIG. **1** is a wearable electronic device **1402**, illustrated as a "smart watch." The wearable electronic device **1402** includes the microphone(s) **116**, the speaker(s) **118**, and a display screen **1404**. Components of the processor(s) **190**, including the communication session manager **140**, are integrated in the wearable electronic device **1402**.

[0152] In a particular example, the wearable electronic device **1402** is operable to automatically apply a speech-based visual indicator to a video component **134** of a communication session **130**. To illustrate, when a user of the wearable electronic device **1402** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document that the user is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can

be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at the display screen **1404**. The wearable electronic device **1402** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0153] In implementations in which the user of the wearable electronic device **1402** is not the audio presenter, such as when the wearable electronic device **1402** receives the presenter's speech, and optionally screen share content, via a media stream, the communication session manager **140** can process the presenter's speech, identify the text or other visual element that the presenter is speaking about, and add the pointer indicator **154** or other visual indicator **152** for display at the display screen **1404** while the presenter's speech is played out via the speaker(s) **118**.

[0154] FIG. **15** is an implementation **1500** in which one of the devices **102** of FIG. **1** is a wireless speaker and voice activated device **1502**. The wireless speaker and voice activated device **1502** can have wireless network connectivity and is configured to execute an assistant operation. The wireless speaker and voice activated device **1502** includes the microphone(s) **116** and the speaker(s) **118**. Components of the processor(s) **190**, including the communication session manager **140**, are integrated in the wireless speaker and voice activated device **1502**.

[0155] In a particular example, the wireless speaker and voice activated device **1502** is operable to automatically apply a speech-based visual indicator to a video component **134** of a communication session **130**. To illustrate, when a user of the wireless speaker and voice activated device **1502** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document that the user is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at a local display, such as a display screen of the wireless speaker and voice activated device **1502** or a display device (e.g., a television) that is coupled to the wireless speaker and voice activated device **1502**. The wireless speaker and voice activated device **1502** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0156] In implementations in which the user of the wireless speaker and voice activated device **1502** is not the audio presenter, such as when the wireless speaker and voice activated device **1502** receives the presenter's speech and screen share content via a media stream, the communication session manager **140** can process the presenter's speech, identify the text or other visual element that the presenter is speaking about, and add the pointer indicator **154** or other visual indicator **152** for display while the presenter's speech is played out via the speaker(s) **118**.

[0157] FIG. **16** depicts an implementation **1600** in which one of the devices **102** of FIG. **1** is a portable electronic device that corresponds to a camera device **1602**. The camera device **1602** includes the microphone(s) **116**, the speaker(s) **118**, and optionally a display screen (e.g., on a side not visible in FIG. **16**). Components of the processor(s) **190**, including the communication session manager **140**, are integrated in the camera device **1602**.

[0158] In a particular example, the camera device **1602** is operable to automatically apply a speech-based visual indicator to a video component **134** of a communication session **130**. To illustrate, when a user of the camera device **1602** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document (e.g., a captured image) that the user is speaking about,

and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at a local display, such as the display screen of the camera device **1602**. The camera device **1602** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0159] FIG. **17** depicts an implementation **1700** in which one of the devices **102** of FIG. **1** is a portable electronic device that corresponds to an extended reality headset **1702** (e.g., a virtual reality, mixed reality, or augmented reality headset). The extended reality headset **1702** includes the microphone(s) **116**, the speaker(s) **118**, and a display device **1704**, such as a display screen disposed on a surface that is positioned in front of a user's eyes when the extended reality headset **1702** is worn. Components of the processor(s) **190**, including the communication session manager **140**, are integrated in the extended reality headset **1702**.

[0160] In a particular example, the extended reality headset **1702** is operable to automatically apply a speech-based visual indicator to a video component **134** of a communication session **130**, such as a virtual conference, seminar, or multi-participant extended reality session (e.g., a multi-player extended reality game). To illustrate, when a user of the extended reality headset **1702** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document or a shared virtual object that includes text (e.g., a map or sign in a VR game) that the user is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at the display device **1704**. For example, the display device **1704** may be coupled to the processor(s) **190**, and the processor(s) **190** may be configured to display the updated video component **136** at the display device **1704**. The extended reality headset **1702** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0161] In implementations in which the user of the extended reality headset **1702** is not the audio presenter, such as when the extended reality headset **1702** receives the presenter's speech and shared content via a media stream, the communication session manager **140** can process the presenter's speech, identify the text or other visual element that the presenter is speaking about, and add the pointer indicator **154** or other visual indicator **152** for display at the display device **1704** while the presenter's speech is played out via the speaker(s) **118**.

[0162] FIG. **18** depicts an implementation **1800** in which one of the devices **102** of FIG. **1** corresponds to, or is integrated within, a vehicle **1802**, illustrated as a manned or unmanned aerial device (e.g., a drone capable of facilitating communication sessions, such as a conference call drone or an air taxi). The vehicle **1802** includes the microphone(s) **116** and the speaker(s) **118**. Components of the processor(s) **190**, including the communication session manager **140**, are integrated in the vehicle **1802**.

[0163] In a particular example, a user of the vehicle **1802** may be a participant in a communication session **130**, and the vehicle **1802** is operable to automatically apply a speech-based visual indicator to a video component **134** of the communication session **130**. For example, the microphone(s) **116** are configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators **152** during the

communication session **130**, such as while the participant is an occupant of the vehicle **1802**. To illustrate, when a user of the vehicle **1802** is an audio presenter during the communication session **130**, the user's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document that the user is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at a display screen (not shown). The vehicle **1802** thus enables the user to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the user's speech.

[0164] In implementations in which the user of the vehicle **1802** is not the audio presenter, such as when the vehicle **1802** receives the presenter's speech and screen share content via a media stream, the communication session manager **140** can process the presenter's speech, identify the text or other visual element that the presenter is speaking about, and add the pointer indicator **154** or other visual indicator **152** for display at a display screen while the presenter's speech is played out via the speaker(s) **118**.

[0165] FIG. **19** depicts another implementation **1900** in which one of the devices **102** of FIG. **1** corresponds to, or is integrated within, a vehicle **1902**, illustrated as a car. The vehicle **1902** includes the processor **190** including the communication session manager **140**. The vehicle **1902** also includes the microphone(s) **116**, the speaker(s) **118**, and a display screen **1904**.

[0166] In a particular example, an occupant of the vehicle **1902** may be a participant in a communication session **130**, and the vehicle **1902** is operable to automatically apply a speech-based visual indicator to a video component **134** of the communication session **130**. For example, the microphone(s) **116** are configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators **152** during the communication session **130**. To illustrate, when an occupant of the vehicle **1902** is an audio presenter during the communication session **130**, the occupant's speech may be captured by the microphone(s) **116** and processed to identify text (or some other element) of a shared document that the occupant is speaking about, and a pointer indicator **154** (or other visual indicator **152**) is added to emphasize or direct a viewer's attention to the identified text or element. The resulting updated video component **136** can be transmitted to one or more other devices associated with the communication session **130**, such as one or more end-user device, communication session server, display device (e.g., a projection device for a class or seminar), or a combination thereof. In addition, the updated video component **136** may also be displayed at the display screen **1904**. The vehicle **1902** thus enables the occupant to conduct a hands-free presentation in which the pointer indicator **154** is automatically added and updated responsive to the content of the occupant's speech.

[0167] In implementations in which the occupant of the vehicle **1902** is not the audio presenter, such as when the vehicle **1902** receives the presenter's speech and screen share content via a media stream, the communication session manager **140** can process the presenter's speech, identify the text or other visual element that the presenter is speaking about, and add the pointer indicator **154** or other visual indicator **152** for display at the display screen **1904** while the presenter's speech is played out via the speaker(s) **118**.

[0168] Referring to FIG. **20**, a particular implementation of a method **2000** of applying a speech-based visual indicator to a video component of a communication session is shown. In a particular aspect, one or more operations of the method **2000** are performed by at least one of the devices **102** of FIG. **1**, the communication session server(s) **106**, the communication session manager **140**, the processor(s) **190**, the system **100**, or a combination thereof.

[0169] The method **2000** includes, at block **2002**, detecting, at a device and during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. For example, the audio analyzer **142** at the device **102A** detects, during the communication session **130**, that the audio component **132** includes particular speech corresponding to the topic **144**.

[0170] The method **2000** includes, at block **2004**, detecting, at the device, that the video component includes an object that is associated with the particular speech. For example, the video analyzer **146**, the speech-to-object matcher **240**, or both, process the video component **134** to detect the object(s) **148**, and selects one of the object(s) **148** (e.g., the selected object **242**) as corresponding to the topic **144**. As another example, the document analyzer **222**, the speech-to-object matcher **240**, or both, process the presentation document(s) **220** to detect the presentation object(s) **224**, and selects one of the presentation object(s) **224** (e.g., the selected object **242**) as corresponding to the topic **144**.

[0171] The method **2000** includes, at block **2006**, updating, at the device, the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting. For example, the video updater **150** updates the video component **134** to generate the updated video component **136** by adding the visual indicator **152**, such as the pointer indicator **154**, the text effect **390**, or the highlighting **392** to the selected object.

[0172] In some implementations, the method **2000** includes, during the communication session: processing the audio component to detect a discussion topic, performing a search of one or more data sources to locate a diagram associated with the discussion topic, and updating the video component to include the diagram. For example, the audio analyzer **142** processes the audio component **132** to determine the topic **144**, which the image retriever **280** sends to the content server(s) **292** as the discussion topic **282** to conduct a search to locate the diagram(s) **284**. The video updater **150** may then update the video component **134** to generate the updated video component **136** that includes the diagram **284**.

[0173] In some implementations, the method **2000** includes, during the communication session, detecting that the audio component includes an utterance that is mapped to a particular action associated with a presentation document, and updating the video component to depict a result of performance of the particular action. For example, the audio analyzer **142** (e.g., the jargon detector **430**) detects an utterance that is sent as the jargon **212** to the jargon mapper **230** (e.g., the jargon database **750**), which maps the utterance to a particular usage or action (e.g., an action **234**). The video updater **150** updates the video component **134** to depict a result of performance of the particular usage or action.

[0174] In some implementations, the method **2000** includes, during the communication session, detecting that the audio component includes a description of a component of a presentation object and, based on determining that the component of the presentation object is not present in the video component, updating the video component to generate a representation of the component of the presentation object based on the description. For example, the audio analyzer **142**, the object generator **270**, or both, may detect that the keywords **210** include the description **1000** of a component **1012** of the presentation object **1010**. Based on a determination that the component **1012** of the presentation object **1010** is not present in the video component **134** (e.g., via comparison to the object(s) **148**), the object generator **270** may cause the video updater **150** to generate a representation of the component **1012**, such as described with reference to the block diagram **1050** of FIG. **10**.

[0175] One technical advantage of the method **2000** is that presenter of the communication session **130** can engage in a hands-free presentation that automatically provides visual emphasis to portions of the displayed slide(s) or page(s) that are currently being discussed. As a result, the presenter's hands are free to engage in other activities, such as holding notecards, gesturing, interacting with a keyboard or pointing device to search notes or perform internet searches, etc., while the

presentation is ongoing, while still providing the other participants with the functionality that is conventionally provided by a manually operated pointer to direct the viewers' attention. Further, the automatic insertion of the visual indicator based on the topic of discussion can provide a more accurate and reliable guide for viewers as compared to conventional manually operated pointers that the presenter may forget to update or which may drift from their intended locations, such as due to physical movement at the presenter's device. In addition, by processing the video component to detect the objects, such as when the video component is received as streaming video from a remote presenter, the video can be automatically enhanced at the receiving device by inserting the visual indicator without requiring access to the source documents being displayed.

[0176] The method **2000** of FIG. **20** may be implemented by a field-programmable gate array (FPGA) device, an application-specific integrated circuit (ASIC), a processing unit such as a central processing unit (CPU), a digital signal processor (DSP), a controller, another hardware device, firmware device, or any combination thereof. As an example, the method **2000** of FIG. **20** may be performed by a processor that executes instructions, such as described with reference to FIG. **21**.

[0177] Referring to FIG. **21**, a block diagram of a particular illustrative implementation of a device is depicted and generally designated **2100**. In various implementations, the device **2100** may have more or fewer components than illustrated in FIG. **21**. In an illustrative implementation, the device **2100** may correspond to, include, or be included within one of the devices **102** of FIG. **1**, or one of the communication session server(s) **106** of FIG. **1**. In an illustrative implementation, the device **2100** may perform one or more operations described with reference to FIGS. **1-20**.

[0178] In a particular implementation, the device **2100** includes a processor **2106** (e.g., a central processing unit (CPU)). The device **2100** may include one or more additional processors **2110** (e.g., one or more DSPs). In a particular aspect, the processor(s) **190** of FIG. **1** correspond to the processor **2106**, the processor(s) **2110**, or a combination thereof. The processor(s) **2110** may include a speech and music coder-decoder (CODEC) **2108** that includes a voice coder ("vocoder") encoder **2136**, a vocoder decoder **2138**, the communication session manager **140**, or a combination thereof.

[0179] The device **2100** may include a memory **2186** and a CODEC **2134**. The memory **2186** may include instructions **2156**, that are executable by the one or more additional processors **2110** (or the processor **2106**) to implement the functionality described with reference to the communication session manager **140**.

[0180] The device **2100** may include a modem **2154** coupled, via a transceiver **2150**, to an antenna **2152**. In implementations in which the device **2100** corresponds to one of the devices **102** of FIG. **1**, the modem **2154** corresponds to the modem **126** of FIG. **1**.

[0181] The device **2100** may include a display **2128** coupled to a display controller **2126**. In implementations in which the device **2100** corresponds to the device **102A** of FIG. **1**, the display **2128** corresponds to the display device **104**, and the speaker(s) **118** and the microphone(s) **116** are coupled to the CODEC **2134**. The CODEC **2134** may include a digital-to-analog converter (DAC) **2102**, an analog-to-digital converter (ADC) **2104**, or both. In a particular implementation, the CODEC **2134** may receive analog signals from the microphone(s) **116**, convert the analog signals to digital signals using the analog-to-digital converter **2104**, and provide the digital signals to the speech and music codec **2108**. The speech and music codec **2108** may process the digital signals, and the digital signals may further be processed by the communication session manager **140**. In a particular implementation, the speech and music codec **2108** may provide digital signals to the CODEC **2134**. The CODEC **2134** may convert the digital signals to analog signals using the digital-to-analog converter **2102** and may provide the analog signals to the speaker(s) **118**.

[0182] In a particular implementation, the device **2100** may be included in a system-in-package or system-on-chip device **2122**. In a particular implementation, the memory **2186**, the processor **2106**, the processor(s) **2110**, the display controller **2126**, the CODEC **2134**, the modem **2154**, and

optionally the transceiver **2150** are included in the system-in-package or system-on-chip device **2122**. In a particular implementation, an input device **2130** and a power supply **2144** are coupled to the system-in-package or the system-on-chip device **2122**. Moreover, in a particular implementation, as illustrated in FIG. **21**, the display **2128**, the input device **2130**, the speaker(s) **118**, the microphone(s) **116**, the antenna **2152**, and the power supply **2144** are external to the system-in-package or the system-on-chip device **2122**. In a particular implementation, each of the display **2128**, the input device **2130**, the speaker(s) **118**, the microphone(s) **116**, the antenna **2152**, and the power supply **2144** may be coupled to a component of the system-in-package or the system-on-chip device **2122**, such as an interface or a controller.

[0183] The device **2100** may include a conference call or video call control device, a smart speaker, a speaker bar, a mobile communication device, a smart phone, a cellular phone, a laptop computer, a computer, a tablet, a personal digital assistant, a display device, a television, a gaming console, a music player, a radio, a digital video player, a digital video disc (DVD) player, a tuner, a camera, a navigation device, a vehicle, a headset, an extended reality headset, an augmented reality headset, a mixed reality headset, a virtual reality headset, an aerial vehicle, a home automation system, a voice-activated device, a wireless speaker and voice activated device, a portable electronic device, a car, a computing device, a communication device, an internet-of-things (IoT) device, a virtual reality (VR) device, a base station, a mobile device, or any combination thereof.

[0184] In conjunction with the described implementations, an apparatus includes means for detecting, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session. For example, the means for detecting that the audio component includes the particular speech can correspond to one of the devices **102**, the processor(s) **190**, the communication session manager **140**, the audio analyzer **142**, the speech-to-text network **202**, the feature extractor **402**, the acoustic model **406**, the decoder **410**, the language model **412**, the topic detector **420**, the integrated circuit **1102** of FIG. **11**, the device **2100** of FIG. **21**, the processor **2106**, the processor(s) **2110**, one or more other circuits or components configured to detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session, or any combination thereof.

[0185] In conjunction with the described implementations, the apparatus includes means for detecting that the video component includes an object that is associated with the particular speech. For example, the means for detecting that the video component includes an object that is associated with the particular speech can correspond to one of the devices **102** of FIG. **1**, the processor(s) **190**, the communication session manager **140**, the video analyzer **146**, the image-to-object network **204**, the document analyzer **222**, the speech-to-object matcher **240**, the pre-processor **510**, the detection model **520**, the integrated circuit **1102** of FIG. **11**, the device **2100** of FIG. **21**, the processor **2106**, the processor(s) **2110**, one or more other circuits or components configured to detect that the video component includes an object that is associated with the particular speech, or any combination thereof.

[0186] In conjunction with the described implementations, the apparatus includes means for updating the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting. For example, the means for updating the video component can correspond to one of the devices **102** of FIG. **1**, the processor(s) **190**, the communication session manager **140**, the video updater **150**, the integrated circuit **1102** of FIG. **11**, the device **2100** of FIG. **21**, the processor **2106**, the processor(s) **2110**, one or more other circuits or components configured to update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting, or any combination thereof.

[0187] In some implementations, a non-transitory computer-readable medium (e.g., a computer-

readable storage device, such as the memory **2186**) includes instructions (e.g., the instructions **2156**) that, when executed by one or more processors (e.g., the processor(s) **190**, the processor(s) **2110**, or the processor **2106**), cause the one or more processors to detect, during a communication session (the communication session **130**) that includes an audio component (e.g., the audio component **132**) and a video component (the video component **134**), that the audio component includes particular speech (e.g., the topic **144**) of a participant (e.g., a person **110**) of the communication session, detect that the video component includes an object (e.g., an object **148**) that is associated with the particular speech, and apply a visual indicator (e.g., the visual indicator **152**) to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

Particular Aspects of the Disclosure are Described Below in Sets of Interrelated Examples:

[0188] According to Example 1, a device comprises: one or more processors configured to: detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session; detect that the video component includes an object that is associated with the particular speech; and update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0189] Example 2 includes the device of Example 1, wherein the one or more processors are configured to use a speech-to-text network to process the audio component to detect the particular speech.

[0190] Example 3 includes the device of Example 1 or Example 2, wherein the one or more processors are configured to use an image-to-object network to process the video component to detect the object.

[0191] Example 4 includes the device of any of Example 1 to Example 3, wherein the object includes at least one of: text, a bullet point, a table element, or a graphical element that includes text.

[0192] Example 5 includes the device of any of Example 1 to Example 3, wherein the object includes text.

[0193] Example 6 includes the device of any of Example 1 to Example 3, wherein the object includes a graphical element that includes text.

[0194] Example 7 includes the device of any of Example 1 to Example 3, wherein the object includes a bullet point.

[0195] Example 8 includes the device of any of Example 1 to Example 3, wherein the object includes a table element.

[0196] Example 9 includes the device of any of Example 1 to Example 8, wherein the particular speech corresponds to a keyword or a topic.

[0197] Example 10 includes the device of any of Example 1 to Example 8, wherein the particular speech corresponds to a keyword.

[0198] Example 11 includes the device of any of Example 1 to Example 8, wherein the particular speech corresponds to or a topic.

[0199] Example 12 includes the device of any of Example 1 to Example 11, wherein the communication session corresponds to at least one of: a conference call, a seminar, or a multi-participant extended reality session.

[0200] Example 13 includes the device of any of Example 1 to Example 11, wherein the communication session corresponds to a conference call.

[0201] Example 14 includes the device of any of Example 1 to Example 11, wherein the communication session corresponds to a seminar.

[0202] Example 15 includes the device of any of Example 1 to Example 11, wherein the communication session corresponds to a multi-participant extended reality session.

[0203] Example 16 includes the device of any of Example 1 to Example 15, wherein the one or

more processors are further configured to, prior to a start of the communication session, analyze one or more presentation documents associated with the communication session to detect presentation objects in the one or more presentation documents, and wherein, during the communication session, detection of the object is at least partially based on a determination that at least one of the presentation objects is associated with the particular speech.

[0204] Example 17 includes the device of any of Example 1 to Example 16, wherein the one or more processors are further configured to: process, during the communication session, the audio component to detect a discussion topic; perform a search of one or more networks to locate a diagram associated with the discussion topic; and update the video component to include the diagram.

[0205] Example 18 includes the device of any of Example 1 to Example 17, wherein the one or more processors are further configured to: detect, during the communication session, that the audio component includes an utterance that is mapped to a particular action associated with a presentation document; and update the video component to depict a result of performance of the particular action.

[0206] Example 19 includes the device of Example 18, wherein the one or more processors are further configured to: for each page, of a presentation document, that was previously displayed during the communication session, track a position of a most recent visual indicator that was depicted on that page; and based on the particular action corresponding to re-displaying a particular page of the previously displayed pages, update the video component to display the particular page and to restore the most recent visual indicator for the particular page.

[0207] Example 20 includes the device of any of Example 1 to Example 19, wherein the one or more processors are further configured to: detect, during the communication session, that the audio component includes a description of a component of a presentation object; and based on determining that the component of the presentation object is not present in the video component, update the video component to generate a representation of the component of the presentation object based on the description.

[0208] Example 21 includes the device of Example 20, wherein the component of the presentation object includes an element of a flow chart, a timeline, or a block diagram.

[0209] Example 22 includes the device of any of Example 1 to Example 21, wherein the one or more processors are included in an end-user device of the participant of the communication session.

[0210] Example 23 includes the device of any of Example 1 to Example 22, wherein the one or more processors are included in an end-user device configured to receive, from a remote device, a media stream that includes the audio component and the video component.

[0211] Example 24 includes the device of any of Example 1 to Example 22, wherein the one or more processors are included in an end-user device configured to send a media stream that includes the audio component and the updated video component to a remote device.

[0212] Example 25 includes the device of any of Example 1 to Example 24, wherein the one or more processors are integrated in a headset that further comprises a microphone coupled to the one or more processors, and wherein the microphone is configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators during the communication session while the participant is engaged in physical activity.

[0213] Example 26 includes the device of Example 25, wherein the headset corresponds to an extended reality headset that further comprises a display device coupled to the one or more processors, and wherein the one or more processors are further configured to display the updated video component at the display device.

[0214] Example 27 includes the device of any of Example 1 to Example 24, wherein the one or more processors are integrated in a vehicle that further comprises a microphone coupled to the one or more processors, and wherein the microphone is configured to capture the particular speech of

the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators during the communication session while the participant is an occupant of the vehicle.

[0215] Example 28 includes the device of any of Example 1 to Example 27, further comprising a speaker configured to play out sound based on the audio component.

[0216] Example 29 includes the device of any of Example 1 to Example 28, further comprising a microphone configured to capture the particular speech of the participant.

[0217] Example 30 includes the device of any of Example 1 to Example 29, further comprising a display device configured to display the updated video component.

[0218] Example 31 includes the device of any of Example 1 to Example 21, wherein the one or more processors are included in a communication session server configured to send the updated video component to one or more end-user devices of participants of the communication session.

[0219] Example 32 includes the device of any of Example 1 to Example 31, further comprising a modem configured to receive at least one of the audio component or the video component from a remote device.

[0220] According to Example 33, a method comprises: detecting, at a device and during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session; detecting, at the device, that the video component includes an object that is associated with the particular speech; and updating, at the device, the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0221] Example 34 includes the method of Example 33, further comprising, during the communication session: processing the audio component to detect a discussion topic; performing a search of one or more data sources to locate a diagram associated with the discussion topic; and updating the video component to include the diagram.

[0222] Example 35 includes the method of Example 33 or Example 34, further comprising, during the communication session: detecting that the audio component includes an utterance that is mapped to a particular action associated with a presentation document; and updating the video component to depict a result of performance of the particular action.

[0223] Example 36 includes the method of any of Example 33 to Example 35, further comprising, during the communication session: detecting that the audio component includes a description of a component of a presentation object; and based on determining that the component of the presentation object is not present in the video component, updating the video component to generate a representation of the component of the presentation object based on the description.

[0224] Example 37 includes the method of Example 36, wherein the component of the presentation object includes an element of a flow chart, a timeline, or a block diagram.

[0225] Example 38 includes the method of any of Example 33 to Example 37, wherein the object includes at least one of: text, a bullet point, a table element, or a graphical element that includes text.

[0226] Example 39 includes the method of any of Example 33 to Example 37, wherein the object includes text.

[0227] Example 40 includes the method of any of Example 33 to Example 37, wherein the object includes a graphical element that includes text.

[0228] Example 41 includes the method of any of Example 33 to Example 37, wherein the object includes a bullet point.

[0229] Example 42 includes the method of any of Example 33 to Example 37, wherein the object includes a table element.

[0230] Example 43 includes the method of any of Example 33 to Example 42, wherein the particular speech corresponds to a keyword or a topic.

[0231] Example 44 includes the method of any of Example 33 to Example 42, wherein the particular speech corresponds to a keyword.

[0232] Example 45 includes the method of any of Example 33 to Example 42, wherein the particular speech corresponds to a topic.

[0233] Example 46 includes the method of any of Example 33 to Example 45, wherein the communication session corresponds to at least one of: a conference call, a seminar, or a multi-participant extended reality session.

[0234] Example 47 includes the method of any of Example 33 to Example 45, wherein the communication session corresponds to a conference call.

[0235] Example 48 includes the method of any of Example 33 to Example 45, wherein the communication session corresponds to a seminar.

[0236] Example 49 includes the method of any of Example 33 to Example 45, wherein the communication session corresponds to a multi-participant extended reality session.

[0237] Example 50 includes the method any of Example 33 to Example 49, further comprising, prior to a start of the communication session, analyzing one or more presentation documents associated with the communication session to detect presentation objects in the one or more presentation documents, and wherein, during the communication session, detection of the object is at least partially based on a determination that at least one of the presentation objects is associated with the particular speech.

[0238] Example 51 includes the method of any of Example 33 to Example 50, wherein the device corresponds to an end-user device of the participant of the communication session.

[0239] Example 52 includes the method of any of Example 33 to Example 50, further comprising receiving, from a remote device, a media stream that includes the audio component and the video component, and wherein the device corresponds to an end-user device.

[0240] Example 53 includes the method of any of Example 33 to Example 50, wherein the device corresponds to an end-user device, and further comprising sending a media stream that includes the audio component and the updated video component to a remote device.

[0241] Example 54 includes the method of any of Example 33 to Example 53, wherein the device corresponds to a headset, and further comprising capturing the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators during the communication session while the participant is engaged in physical activity.

[0242] Example 55 includes the method of Example 54, wherein the headset corresponds to an extended reality headset, and further comprising displaying the updated video component at a display device of the extended reality headset.

[0243] Example 56 includes the method of any of Example 33 to Example 53, wherein the device corresponds to a vehicle, and further comprising capturing the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators during the communication session while the participant is an occupant of the vehicle.

[0244] Example 57 includes the method of any of Example 33 to Example 56, further comprising playing out, via a speaker, sound based on the audio component.

[0245] Example 58 includes the method of any of Example 33 to Example 57, further comprising capturing, via a microphone, the particular speech of the participant.

[0246] Example 59 includes the method of any of Example 33 to Example 59, further comprising displaying the updated video component at a display device.

[0247] According to Example 60, a device comprises: a memory configured to store instructions; and a processor configured to execute the instructions to perform the method of any of Example 33 to Example 59.

[0248] According to Example 61, a non-transitory computer-readable medium stores instructions that, when executed by one or more processors, cause the one or more processors to perform the method of any of Example 33 to Example 59.

[0249] According to Example 62, an apparatus comprising means for performing the method of any of Example 33 to Example 59.

[0250] According to Example 63, a non-transitory computer-readable medium stores instructions that, when executed by one or more processors, cause the one or more processors to: detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session; detect that the video component includes an object that is associated with the particular speech; and update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0251] According to Example 64, an apparatus comprises: means for detecting, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session; means for detecting that the video component includes an object that is associated with the particular speech; and means for updating the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting.

[0252] Those of skill would further appreciate that the various illustrative logical blocks, configurations, modules, circuits, and algorithm steps described in connection with the implementations disclosed herein may be implemented as electronic hardware, computer software executed by a processor, or combinations of both. Various illustrative components, blocks, configurations, modules, circuits, and steps have been described above generally in terms of their functionality. Whether such functionality is implemented as hardware or processor executable instructions depends upon the particular application and design constraints imposed on the overall system. Skilled artisans may implement the described functionality in varying ways for each particular application, such implementation decisions are not to be interpreted as causing a departure from the scope of the present disclosure.

[0253] The steps of a method or algorithm described in connection with the implementations disclosed herein may be embodied directly in hardware, in a software module executed by a processor, or in a combination of the two. A software module may reside in random access memory (RAM), flash memory, read-only memory (ROM), programmable read-only memory (PROM), erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), registers, hard disk, a removable disk, a compact disc read-only memory (CD-ROM), or any other form of non-transient storage medium known in the art. An exemplary storage medium is coupled to the processor such that the processor may read information from, and write information to, the storage medium. In the alternative, the storage medium may be integral to the processor. The processor and the storage medium may reside in an application-specific integrated circuit (ASIC). The ASIC may reside in a computing device or a user terminal. In the alternative, the processor and the storage medium may reside as discrete components in a computing device or user terminal.

[0254] The previous description of the disclosed aspects is provided to enable a person skilled in the art to make or use the disclosed aspects. Various modifications to these aspects will be readily apparent to those skilled in the art, and the principles defined herein may be applied to other aspects without departing from the scope of the disclosure. Thus, the present disclosure is not intended to be limited to the aspects shown herein but is to be accorded the widest scope possible consistent with the principles and novel features as defined by the following claims.

Claims

1. A device comprising: one or more processors configured to: detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session; detect that the video component includes an object that is associated with the particular speech; update the video component to apply a visual indicator to the object, the visual indicator including at least one of a

pointer indicator, a text effect, or highlighting; and track a most recent position of the visual indicator to enable the visual indicator to be restored to the most recent position after a subsequent update of the video component.

2. The device of claim 1, wherein, to track the most recent position of the visual indicator, the one or more processors are configured to maintain a pointer position history, and wherein the pointer position history indicates the most recent position of the visual indicator.

3. The device of claim 2, wherein the pointer position history includes a data structure that has a plurality of entries, each entry of which associates the video component with a position of the visual indicator.

4. The device of claim 1, wherein the one or more processors are configured to: use a speech-to-text network to process the audio component to detect the particular speech; and use an image-to-object network to process the video component to detect the object.

5. The device of claim 4, wherein the image-to-object network includes a detection model, wherein the detection model includes a plurality of classifiers, and wherein a classifier of the plurality of classifiers is configured to detect an object type associated with an image of the video component.

6. The device of claim 1, wherein the one or more processors are configured to: process, during the communication session, the audio component to detect a discussion topic; perform a search of one or more networks to locate a diagram associated with the discussion topic; and update the video component to include the diagram.

7. The device of claim 1, wherein the one or more processors are configured to: detect, during the communication session, that the audio component includes an utterance that is mapped to a particular action associated with the video component; and update the video component to depict a result of performance of the particular action.

8. The device of claim 1, wherein the one or more processors are further configured to restore the visual indicator to the most recent position after the subsequent update of the video component.

9. The device of claim 1, wherein the one or more processors are further configured to: detect, during the communication session, that the audio component includes a description of a component of a presentation object; and based on determining that the component of the presentation object is not present in the video component, update the video component to generate a representation of the component of the presentation object based on the description.

10. The device of claim 1, wherein the one or more processors are integrated in a headset that further comprises a microphone coupled to the one or more processors, and wherein the microphone is configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators during the communication session while the participant is engaged in physical activity.

11. The device of claim 1, wherein the one or more processors are integrated in a vehicle that further comprises a microphone coupled to the one or more processors, and wherein the microphone is configured to capture the particular speech of the participant to enable the participant to deliver a hands-free presentation with automatic visual indicators during the communication session while the participant is an occupant of the vehicle.

12. The device of claim 1, further comprising a modem configured to receive at least one of the audio component or the video component from a remote device.

13. The device of claim 1, further comprising a speaker configured to play out sound based on the audio component.

14. The device of claim 1, further comprising a microphone configured to capture the particular speech of the participant.

15. The device of claim 1, further comprising a display device configured to display the updated video component.

16. A method comprising: detecting, at a device and during a communication session that includes an audio component and a video component, that the audio component includes particular speech

of a participant of the communication session; detecting, at the device, that the video component includes an object that is associated with the particular speech; updating, at the device, the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting; and tracking a most recent position of the visual indicator to enable the visual indicator to be restored to the most recent position after a subsequent update of the video component.

17. The method of claim 16, wherein: tracking the most recent position of the visual indicator includes maintaining a pointer position history; the pointer position history indicates the most recent position of the visual indicator; and the pointer position history includes a data structure that has a plurality of entries, each entry of which associates the video component with a position of the visual indicator.

18. The method of claim 16, wherein detecting, at the device, that the video component includes an object that is associated with the particular speech includes: converting, by a pre-processor of the device, a frame of the video component into an enhanced greyscale image; converting, by the pre-processor, the enhanced greyscale image into a binary image; generating, by a detection model of the device, a plurality of feature maps associated with features of the binary image; identifying, by the detection model and based on the plurality of feature maps, a segment of the binary image that includes text; and detecting, by the detection model, an object associated with the segment.

19. A non-transitory computer-readable medium storing instructions that, when executed by one or more processors, cause the one or more processors to: detect, during a communication session that includes an audio component and a video component, that the audio component includes particular speech of a participant of the communication session; detect that the video component includes an object that is associated with the particular speech; update the video component to apply a visual indicator to the object, the visual indicator including at least one of a pointer indicator, a text effect, or highlighting; and track a most recent position of the visual indicator to enable the visual indicator to be restored to the most recent position after a subsequent update of the video component.

20. The non-transitory computer-readable medium of claim 19, wherein: to track the most recent position of the visual indicator, the instructions, when executed by the one or more processors, cause the one or more processors to maintain a pointer position history; the pointer position history indicates the most recent position of the visual indicator; and the pointer position history includes a data structure that has a plurality of entries, each entry of which associates the video component with a position of the visual indicator.
