

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250260835

Kind Code

A1

Publication Date

August 14, 2025

Inventor(s)

BOSSE; Sebastian et al.

MULTI-VIEW CODING WITH EXPLOITATION OF RENDERABLE PORTIONS

Abstract

Rendered portions of a renderable portion of a first view of a multi-view signal are introduced into a prediction loop of a multi-view encoder to form a reference signal for the block-based prediction of the encoder's view predictor may improve the prediction and thereby increasing the coding efficiency. The introduction may be performed by completely inserting the renderable portion into the prediction loop to form new reference pictures with their own reference picture indices in addition to reference pictures obtained by the reconstructed version of the multi-view signal of the block-based prediction. Alternatively, the rendered portion may be introduced into the prediction loop completely, but with replacing the normal prediction reference signal, i.e., the normally reconstructed signal as obtained by block-based prediction. Alternatively, it may be signaled within the data stream which portions of the renderable portion are used to replace respective portions in the normally reconstructed signal.

Inventors: BOSSE; Sebastian (Berlin, DE), SCHWARZ; Heiko (Panketal, DE), WIEGAND; Thomas (Berlin, DE), BRUST; Heribert (Berlin, DE)

Applicant: DOLBY VIDEO COMPRESSION, LLC (Wilmington, DE)

Family ID: 47257758

Appl. No.: 19/051780

Filed: February 12, 2025

Related U.S. Application Data

parent US continuation 18207473 20230608 parent-grant-document US 12244851 child US 19051780

parent US continuation 17129457 20201221 parent-grant-document US 11689738 child US 18207473

parent US continuation 16291146 20190304 parent-grant-document US 10887617 child US 17129457

parent US continuation 14273701 20140509 parent-grant-document US 10264277 child US 16291146
parent WO continuation PCT/EP2012/072180 20121108 PENDING child US 14273701
us-provisional-application US 61558643 20111111

Publication Classification

Int. Cl.: H04N19/52 (20140101); H04N19/597 (20140101)

U.S. Cl.:

CPC H04N19/52 (20141101); H04N19/597 (20141101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a Continuation of U.S. patent application Ser. No. 18/207,473 filed Jun. 8, 2023, which is a continuation of U.S. patent application Ser. No. 17/129,457 filed Dec. 21, 2020 now U.S. Pat. No. 11,689,738, which is a continuation of U.S. patent application Ser. No. 16/921,146 filed Mar. 4, 2019, now U.S. Pat. No. 10,887,617 issued Jan. 5, 2021, which is a continuation of U.S. patent application Ser. No. 14/273,701 filed May 9, 2014 now U.S. Pat. No. 10,264,277 issued Apr. 16, 2019, which is a Continuation of International Application No. PCT/EP2012/072180, filed Nov. 8, 2012, which claims priority from U.S. Application No. 61/558,643, filed Nov. 11, 2011. The subject matter of each of the foregoing applications is incorporated herein by reference in entirety.

BACKGROUND OF THE INVENTION

[0002] The present application is concerned with multi-view coding.

[0003] Multi-view video sequences are basically captured as multiple single view sequences. These single view sequences are captured by multiple cameras simultaneously from different view-points of the same scene. Therefore, multi-view video sequences contain a high amount of inter-view redundancies.

[0004] A common technique to deal with these inter-view redundancies is interview prediction, analogous to the well known temporal motion-compensated or inter-frame prediction. In interview prediction, the reference frame does not relate temporally but spatially (regarding the camera position) to the frame to be coded. Since these two kinds of prediction are conceptually the same, they can be easily combined by using the same reference lists for both kinds of prediction (i.e., a reference picture list can contain both temporal reference pictures as well as inter-view reference pictures).

[0005] Such a combination of temporal and inter-view prediction is used by the MVC extension to H.264/AVC.

[0006] An example for an effective prediction structure combining temporal and inter-view prediction is presented in FIG. 1. On the left side, a possible prediction structure is shown for the 3 view case, on the right side, an example for the 2 view case is given. In both cases, view Vo is the reference view that is used for interview prediction.

[0007] Interview prediction, as used in MVC is a feasible technique to deal with interview redundancies if only few views are transmitted, e.g. in stereoscopic (two-view) video. The transmitted amount of data in MVC increases approximately linearly with the number of views. This makes MVC unsuitable for applications that demand a higher number of views, such as autostereoscopic displays, where 28 or more views are presented. In such a scenario, not all the

views are transmitted, but only a few views, e.g. 3 views. The bigger part of the views is rendered at the decoder side using the transmitted views. In order to decrease complexity of rendering, new approaches in multi-view coding do not only encode texture (as in MVC), but also depth information in form of depth maps plus camera parameters. This provides the receiver with 3D scene information and eases the interpolation (rendering) of intermediate views.

[0008] Due to disocclusions and pixel displacements that are reaching out of the image plane, not all regions of a frame can be rendered from another view.

[0009] FIG. 2 sketches the rendering for a scene that just contains a square **10** in front of a white background (**12**). The right view and the left view are transmitted, the intermediate view is rendered. The regions marked hatched cannot be rendered from the right view, due to disocclusions (cross hatched with dashed lines) and pixel displacements reaching out of the image plane (cross hatched with continuous lines), while analogously the regions marked simply hatched cannot be rendered from the left view. The regions marked white, i.e. the background, and the object **10** in the rendered view are present in both (left and right) transmitted views.

[0010] FIG. 3 shows an example of the rendering process from a left view to a right view. The pixels of the image regions that cannot be rendered are set to black as shown at **14**.

[0011] It becomes obvious that the transmitted views, i.e. left and right views in FIGS. 2 and 3, have almost the same content for a huge part of the image. Since the depth information **16** and camera parameters **18** are usually transmitted anyway in order to support the rendering at the decoder side, regions **14** that can be rendered by renderer **20** from one transmitted view, such as the left view in FIG. 3, to another transmitted view, such as the right view in FIG. 3, only need to be transmitted in the bitstream once. Thus, conceptually, if picture regions that can efficiently be rendered from one view to another, all regions of the right view except regions **14**, are only transmitted once, a significant amount of the overall bit rate can be saved.

[0012] However, even the non-transmission of renderable portions of secondary transmitted views being renderable from primary transmitted views does not lead to an optimally efficient compression of the multi-view data. Accordingly, it would be favorable to have a multi-view concept at hand which enables a more efficient coding of transmitted views.

SUMMARY

[0013] According to an embodiment, an encoder for encoding a multi-view signal into a data stream may have: a determinator configured to determine a renderable portion of a first view of the multi-view signal which is renderable, by view synthesis, from a second view of the multi-view signal; a view predictor configured to perform a block-based prediction of the multi-view signal from a reference signal so as to obtain a prediction signal; a residual generator configured to determine a residual signal between the prediction signal and the multi-view signal view; a data stream generator configured to insert the residual signal into the data stream, wherein the determinator is configured to actually render at least a reference portion of the renderable portion of the first view from the second view so as to obtain a rendered reference signal and the view predictor is configured to perform, the block-based prediction from the reference signal including the rendered reference signal.

[0014] According to another embodiment, a decoder for reconstructing a multi-view signal from a data stream may have: a renderer configured to render, by view synthesis, at least a reference portion of a first view of the multi-view signal from a reconstructed version of the second view so as to obtain a rendered reference signal; a view predictor configured to perform a block-based prediction of the multi-view signal from a reference signal based on prediction parameters within the data stream so as to obtain a prediction signal, a reconstructor configured to reconstruct the multi-view signal based on a residual signal within the data stream and the prediction signal, wherein the view predictor is configured to perform, the block-based prediction from the reference signal including the rendered reference signal.

[0015] According to another embodiment, a method for encoding a multi-view signal into a data

stream may have the steps of: determining a renderable portion of a first view of the multi-view signal which is renderable, by view synthesis, from a second view of the multi-view signal; performing a block-based prediction of the multi-view signal from a reference signal so as to obtain a prediction signal; determining a residual signal between the prediction signal and the multi-view signal view; inserting the residual signal into the data stream, wherein the determining includes actually rendering at least a reference portion of the renderable portion of the first view from the second view so as to obtain a rendered reference signal and the performance of the prediction includes performing the block-based prediction from the reference signal including the rendered reference signal.

[0016] According to another embodiment, a method for reconstructing a multi-view signal from a data stream may have the steps of: rendering, by view synthesis, at least a reference portion of a first view of the multi-view signal from a reconstructed version of the second view so as to obtain a rendered reference signal; performing a block-based prediction of the multi-view signal from a reference signal based on prediction parameters within the data stream so as to obtain a prediction signal, reconstructing the multi-view signal based on a residual signal within the data stream and the prediction signal, wherein the performance of the prediction includes performing the block-based prediction from the reference signal including the rendered reference signal.

[0017] Another embodiment may have a computer program having a program code for performing, when running on a computer, an inventive method.

[0018] The basic idea underlying the present invention is that introducing actually rendered portions of a renderable portion of a first (or dependent) view of a multi-view signal into a prediction loop of a multi-view encoder so as to form a reference signal for the block-based prediction of the encoder's view predictor may improve the prediction and accordingly, increase the coding efficiency. The introduction may be performed in several manners, such as completely inserting the renderable portion into the prediction loop so as to form new reference pictures with their own reference picture indices in addition to reference pictures obtained by the reconstructed version of the multi-view signal of the block-based prediction. Alternatively, the actually rendered portion may be introduced into the prediction loop completely, but with replacing the normal prediction reference signal, i.e. the normally reconstructed signal as obtained by block-based prediction. In another alternative, it may be signaled within the data stream which portions of the renderable portion are used to replace respective portions in the normally reconstructed signal.

[0019] This concept is, in accordance with further embodiments of the present invention, combined with another discovery made by the inventors, namely that simply disregarding a renderable portion of a first (or dependent) view of a multi-view signal, which is renderable by synthesis from a second (or reference) view of the multi-view signal, introduces new coding penalties due to the lack of information on prediction parameters relating to the renderable portion in the renderable portion, which could serve as a prediction basis for predicting prediction parameters relating to the non-renderable portion of the multi-view signal, i.e. a remaining portion of the multi-view signal disjoint to the renderable portion of the first view. Particularly at the spatio-temporal neighborhood of the renderable portion of the first view, the simple omission of any prediction parameters relating to the renderable portion of the first view would frequently result in a lack of adequate parameter predictors for coding blocks within this neighborhood, both spatially and temporally as well as in inter-view direction as far as further subsequent views in dependency direction are concerned. Accordingly, the prediction parameters of these blocks would have to be transmitted completely anew, thereby causing coding penalties which, to a large extent, compensate the rate savings obtained by omitting the transmission of any data relating to the renderable portion of the first view. Accordingly, as an optional feature of some embodiments of the present invention an intermediate way of handling the renderable portion of the first view results in a more efficient coding. In particular, instead of omitting the coding of the renderable portion of the first view completely, an even more efficient coding of multi-view signals may be obtained if merely the

coding of the residual signal is suppressed within the renderable portion, whereas the prediction parameter coding continues to take place from the non-renderable portion of the multi-view signal across the renderable portion so that prediction parameters relating to the renderable portion may be exploited for predicting prediction parameters relating to the non-renderable portion of the multi-view signal. The additional coding rate which has to be spent for transmitting the prediction parameters relating to the renderable portion may be kept low as the aim of the transmission of these prediction parameters is merely to form a continuation of the prediction parameter history across the renderable portion so as to serve as a basis for prediction parameters of other portions of the multi-view signal. In other words, the prediction parameters relating to the renderable portion do not have to perfectly predict the texture within the renderable portion of the first view so as to keep the residual signal within the renderable portion low.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0020] Embodiments of the present invention will be detailed subsequently referring to the appended drawings, in which:

[0021] FIG. 1 shows prediction structures of a data stream coded according to the MVC extension to H.264/AVC using intra prediction, temporal prediction (vertical errors) and inter-view prediction (horizontal errors);

[0022] FIG. 2 shows exemplarily the rendition of an intermediate view (rendered view) from left and right views, respectively;

[0023] FIG. 3 shows the rendering of a transmitted dependent view (right view) from a transmitted reference view (left view);

[0024] FIG. 4 shows a block diagram of an encoder for encoding a multi-view signal in accordance with an embodiment;

[0025] FIG. 5 shows a block diagram of a decoder for reconstructing a multi-view signal in accordance with an embodiment; and

[0026] FIG. 6 shows a schematic diagram showing a block to be coded in an un-subdivided state of a hierarchy level of a recursive quad-tree subdivision, and a subdivided state according to a subsequent hierarchy level thereof.

DETAILED DESCRIPTION OF THE INVENTION

[0027] Before describing several embodiments of the present invention with regard to the figures, it should be noted that the same reference signs have been used within these figures for equal elements or elements assuming the same functionality, and accordingly the description of such an element with regard to one figure shall equally apply to another figure.

[0028] FIG. 4 shows an encoder for encoding a multi-view signal **20** into a data stream **21**. As will become clear from the following description, the multi-view signal **20** is exemplarily shown to comprise a first view **22** and a second view **24**. However, the embodiments of the present invention are not restricted to such a two-view case and are equally applicable to a case where the number of views of the multi-view signal **20** is higher than two. In any case, the multi-view signal **20** comprises views **22** and **24** which belong to, or relate to, the same scene, but to different view points, and accordingly a high amount of redundancy exists between views **22** and **24** of multi-view signal **20**. The encoder of FIG. 4 efficiently removes these redundancies from data stream **21**.

[0029] Moreover, as will also become clear from the description below, the encoder of FIG. 4 does not necessarily completely encode multi-view signal **20** into data stream **22**. Rather, the following description concentrates on the encoding of view **24** into data stream **22** although it should be clear that encoder **26** could generally be configured to encode views **22** and **24** of multi-view signal **20** into data stream **22**. Similarly, FIG. 4 exemplarily shows view **22** as comprising a depth/disparity

map so as to alleviate the rendering of view **24** from view **22**. That is, view **22** is shown to comprise spatially sampled texture information **30**, and associated therewith a depth/disparity map **28** representing a spatial sampling of the associated depth/disparity of the scene captured from view **22**. It should be noted that view **22** does not have to have a depth/disparity map **28**, as will be outlined in more detail below, wherein, however, the existence of the depth/disparity map **28** is advantageous due to the alleviation of the rendering process at the decoder side as will be outlined in more detail below.

[0030] As far as the dependent view **24** is concerned, same comprises a spatial sampling **32** of texture of view **24**. Merely as an optional feature, view **24** may also comprise depth/disparity map data **34** for associating depth/disparity values to the texture samples of texture **32**, just as depth/disparity map **28** does with respect to texture **30**.

[0031] Further, the multi-view signal **20** may, as shown in FIG. **4**, not only represent a spatial sampling of one common scene from different view directions associated with views **22** and **24**, but also a temporal sampling of this scene. Accordingly, the texture information **30** and **32** may be videos, i.e. sequences of pictures, and the depth/disparity maps **28** and **34** may have a depth/disparity sample array associated with each picture of videos **30** and **32**, respectively. However, although FIG. **4** suggests that a depth/disparity map exists for each picture of the respective view, this is not necessarily the case. In other words, the spatial and/or temporal resolution of the depth/disparity maps **28** and **34**, respectively, may differ from the spatial and/or temporal sampling resolution of the associated texture information **30** and **32**, respectively. The same applies with respect to the spatial and/or temporal resolution of views **22** and **24**: same may be the same among views **22** and **24** or different.

[0032] As is shown in FIG. **4**, encoder **26** comprises a determinator **36**, a view predictor **38**, a residual generator **40**, a parameter predictor **42** and a data stream generator **44**. The interconnection of these blocks is as indicated in FIG. **4** with solid lines. Dashed interconnections and blocks are optional with their significance being derivable from the below-outlined explanation. Regarding the determinator block it is noted that same occurs twice, with position **92** however, representing an alternative position to block **36** so that block **36** may absent when the determinator functionality is included into block **92**, for example.

[0033] The determinator **36** is configured to determine a renderable portion **46** of view **24**, which is renderable by view synthesis from view **22** or any other view belonging to the same scene of the multi-view signal. The rendition by view synthesis, which may also called warping, may, for example, involve shifting each sample in view **22**, such as each picture sample of video **30** of view **22**, in accordance with the co-located value of the depth/disparity map so as to obtain this scene as seen from the view direction of view **24**, followed by interpolating the thus obtained rendition result onto the sample grid of view **24**, such as the sample positions of video **32** of view **24**.

[0034] It should be emphasized that determinator **36** may not actually perform the entire rendition. Rather, it is sufficient for determinator **36** to segment view **24** into the renderable portion **46**, and the remaining portion, i.e. the non-renderable portion **48**. As described above with respect to FIG. **2**, the non-renderable portion **48** results from disocclusions and pixel displacements reaching out of the field of view of view **24**. For example, usually the disparity displacements of the samples of video **30** do not have a vertical component, but are restricted to a horizontal direction. Accordingly, determinator **36** could be configured to row-wise detect sample positions of view **24**, for which no contributing rendered sample from view **22** exists. The actual interpolation would not have to be performed. Even alternatively, determinator **36** could even be configured to detect portions within the depth/disparity map **28** of view **22**, which are likely to cause disocclusions due to a transition from foreground to background in the rendering direction **50** from a view point of view **22** to a view point of view **24**.

[0035] The view predictor **38** is configured to perform a block-based prediction of the multi-view signal from a previously encoded portion of the multi-view signal so as to obtain a prediction

signal **52**. For example, view predictor **38** subdivides each picture of video **32** and **30**, respectively, into blocks and associates respective prediction parameters to these blocks which define how to predict the respective block from a respective previously encoded portion of the multi-view signal **20**. The subdivision of pictures of video **32** may differ from the subdivision of pictures of video **30**. The subdivision may even vary in time using, for example, a quad-tree subdivision of a regular subdivision of the pictures into tree root blocks. The respective subdivision information is, for example, signaled within data stream **22** via data stream generator **44**. The sort of prediction parameters for the blocks may depend on the prediction mode associated with the respective block. For example, the prediction modes may encompass one or more intra prediction modes according to which the respective block is spatially predicted from previously encoded samples neighboring the respective block, one or more inter or temporal prediction modes according to which the respective block is temporally predicted from previously encoded pictures of the same view, and/or one or more inter-view prediction modes according to which the respective block is predicted from a time-aligned picture of another view. In particular, motion compensated prediction may be used for the inter predicted blocks, and disparity compensated prediction may be used for the inter-view predicted blocks. Accordingly, the prediction parameters for inter-predicted blocks may comprise a motion vector indicating the location from where the respective block is to be copied from the reference picture of the same view. A reference index may belong to the motion prediction parameters so as to identify the reference picture among the previously encoded pictures of the same view. Similarly, disparity prediction parameters may comprise a disparity vector and, optionally, a view index.

[0036] As far as the coding order used by view predictor **38** is concerned, same may be chosen such that view **22** is coded prior to view **24** as far as one time instant is concerned. Moreover, the coding order may be such that the coding of the multi-view signal **20** steps to the next time instant merely after having completed the coding of the previously encoded time instant. Moreover, the coding order may be such that the time instants are coded into data stream **22** in the presentation time order or in a different order. Finally, it is noted that view predictor **38** could also be responsible for performing the block-based prediction onto the depth/disparity maps **28** and **34**, respectively. According to the just-mentioned coding order, it could be that video **30** is coded prior to depth/disparity map **28**, followed by view **24**. Similar coding modes as described above for the videos **30** and **32** could be available for depth/disparity maps **28** and **34**, respectively. Additional coding modes could, however, also be available for blocks of depth/disparity maps **28** and **34** with such coding modes, for example, exploiting the similarity and the location of edges in texture on the one hand and depth information on the other hand.

[0037] Although encoder **26** could theoretically use the original version of the previously encoded portion of the multi-view signal as a basis for the prediction, it is advantageous that view predictor **38** bases the prediction on a reference signal **59** which is derived from, or at least partially formed by, a reconstructible version **58** of the respective previously encoded portion of the multi-view signal as it is depicted in FIG. **4**, since in that case a respective view predictor at the decoder side is able to completely recover the prediction signal **52** so that a mismatch between the prediction on the encoder side and the decoder side may be avoided.

[0038] The residual generator **40** is configured to determine a residual signal **54** between the prediction signal **52** and the multi-view signal **20**. The data stream generator **44** may insert the residual signal **54** into data stream **22**. For example, transform coding and/or entropy coding may be used by data stream generator **44**.

[0039] In order to further reduce the coding rate, parameter predictor **42** optionally predicts prediction parameters based on other prediction parameters. To be more precise, the prediction parameters of a current block may be predicted by a parameter predictor **42** based on the prediction parameters of previously encoded blocks of the multi-view signal **20**. For example, parameter predictor **42** predicts the prediction parameters of a current block of video **32** based on prediction

parameters of a previously encoded spatially neighboring block of the same picture or a co-located block of a previously coded picture of video 32.

[0040] The usage of the reconstructed version of the previously encoded portion of the multi-view signal 20 instead of the original version thereof is shown in FIG. 4 by encoder 26 optionally comprising a reconstructor 56, which recovers the reconstructible version 58 of the previously encoded portion of the multi-view signal 20, as reconstructible at the decoder side from the prediction parameters 52 and the residual signal 54 as transmitted via data stream 22.

[0041] In order to effectively encode the multi-view signal 20, the encoder 26 may be configured such that the view predictor 38 continues to perform the prediction even across the renderable portion 46, with the data stream generator 44 inserting the parameter prediction residual 60 into data stream 22, while the residual coder, however, could be configured to skip residual generation within the renderable portion 16. By this measure, the parameter predictor 42 could be able to effectively perform the parameter prediction even across the boundary between the non-renderable portion 48 and the renderable portion 46 of view 24, while spending coding rate for residual coding within the renderable portion 1s avoided. However, it is emphasized that this is an optional feature of FIG. 4. For example, view predictor 38 may perform the block-based prediction using first prediction parameters within the renderable portion 46, and using second prediction parameters within the non-renderable portion, and the parameter predictor 42 may be configured to predict the second prediction parameters from the first prediction parameters so as to obtain respective parameter prediction residual 60 for insertion into the data stream by data stream generator 44. By this measure the coding rate for coding the second prediction parameters could be reduced relative to the case where no first prediction parameters would be available for the decoder. Rather, information of the first prediction parameters, i.e. the prediction parameters used by predictor 38 within the renderable portion 46, would be transmitted via bitstream 22 such as, for example, by predictive coding using parameter predictor 42. In any case, the data stream 22 would have information on the first prediction parameters used within the renderable portion 46 so that, at the decoder side, these prediction parameters may be recovered for the renderable portion 46.

[0042] However, it could also be possible that blocks within the renderable portion 46 are neither predicted by view predictor 38 nor subject to residual coding in residual generator 40 so that both information entities, prediction parameters and the residual signal, would be missing for such blocks within the renderable portion 46. The missing information in the prediction loop between the output of reconstructor 56 and reference input of view predictor 38 could, however, be filled by the rendition of these missing portions based on an actual rendition 57 thereof by renderer 92 as it will be described in more detail below with respect to the various embodiments of the present application. Briefly spoken, actually rendered portions of the renderable portion 46, such as the whole portion 46, may be input used to contribute to the reference signal, such as in form of a whole reference picture. Alternatively, portions of the renderable portion 46, may replace collocated portions within the reconstructed signal 58, with the replacement being performed by a replacer, thereby obtaining a mixed signal 91 having blocks rendered, and block reconstructed.

[0043] However, it could also be possible that blocks within the renderable portion 46 are predicted by view predictor 38 based on the reference signal 59 and that residual coding in residual generator 40 does also take place so that both information entities, prediction parameters and the residual signal, would be present for such blocks within the renderable portion 46. The prediction parameters could be restricted to a pure indication of a coding mode of the respective block, such as a coding mode according to which the co-located portion of the actually rendered signal 57 is copied into the current block within portion 46 as a predictor, also with, optionally, a residual signal or without.

[0044] For example, the view predictor 38 may be configured to, in performing the block-based prediction, predict at least one first block within the renderable portion 46 using first motion or disparity vector data, and at least one second block within the non-renderable portion 48 using

second motion or disparity vector data, and the parameter predictor **42** would be configured to predict the second motion or disparity vector data from the first motion or disparity vector data. In this regard, the first block and the second block may both belong to the same picture of the first view **24** or may both belong to the same time instant, or the first block and the second block may be within different pictures of the first view **24**. However, it could also be that the second block is a block of a third view of the multi-view signal **20**, such as a block belonging to the same time instant as the first block of the renderable portion **46** of view **24**, with the third view, however, following in view coding order defined among the views of multi-view signal **20**.

[0045] As is shown in FIG. **4**, reconstructor **56** is configured to reconstruct the first view **24**, including the renderable portion **46**, from the prediction signal **52** and the residual signal **54** (as far as present), and the view predictor **38** performs the block-based prediction from the reference signal **59** which, at least partially, is based on the resulting reconstructed signal **58**. At least a part of the reference signal **59** is composed of the actually rendered portion **57** of the rendering result of the renderable portion **46**, i.e. the rendered reference signal **57**, which is also be introduced into the prediction loop as will be outlined in more detail below.

[0046] Before this, however, an embodiment of a decoder is described in more detail below, the decoder being able to decode data stream **21**. As will become clear from this description, any lack of the residual signal within data stream **22** as far as the renderable portion **46** is concerned, does not negatively influence the quality of the reproduced view **24** at the decoder side as the decoder is able to render this portion **46** from the reconstructed version of view **22**, which in turn is coded into data stream **22** along with the residual signal **54**, thereby resulting in an appropriate quality of the view synthesis in view **24**.

[0047] FIG. **5** shows such a decoder in accordance with an embodiment. The decoder of FIG. **5** is configured to reconstruct the first view **24** of the multi-view signal **20** from the data stream **22**. Again, the decoder of FIG. **5** could also be configured to reconstruct view **22**, but FIG. **5** concentrates on view

[0048] The decoder of FIG. **5** comprises a data stream extractor **80**, a view predictor **82**, a reconstructor **84**, a renderer **86** and a replacer **88**. The interconnection of these blocks is as indicated in FIG. **5** with solid lines. Dashed interconnections are optional with their significance being derivable from the below-outlined explanation.

[0049] The data stream extractor **80** is configured to extract prediction parameters **53** and a residual signal **54** from data stream **22**. The extraction of the prediction parameters **53** may involve predictive decoding using previously extracted or previously decoded prediction parameters. In other words, data stream extractor **80** may also incorporate a functionality which recovers the prediction parameters **53** from the prediction parameter residuals **60** within data stream **22**, namely by predicting the prediction parameters from previously prediction parameters and refining the predicted prediction parameters by use of the prediction parameter residual **60**.

[0050] The functionality of the view predictor **82** and the reconstructor **84** is to emulate the operation of view predictor **38** and reconstructor **46** of encoder **26** of FIG. **4**. That is, view predictor **82** performs a block-based prediction of the multi-view signal from a previously reconstructed portion **58** of the multi-view signal so as to obtain the prediction signal **52** and the reconstructor **84** is configured to reconstruct the multi-view signal **22** using the prediction signal **52** and the residual signal **54** so as to obtain a reconstructed version of the multi-view signal as obtained without rendering. A presentation of the reconstructed version **58** would could result in a poor quality of the multi-view signal within the renderable portion **46** in case the above-mentioned optional residual suppression within the renderable portion **46** is used according to which the residual signal or both prediction signal and residual signal are, at least partially, missing within the renderable portion **46**. Accordingly, renderer **86** is configured to actually perform the rendering of the renderable portion **46** of the first view **24** from the reconstructed version of view **22**, at least partially, so as to obtain an actually rendered portion **57**.

[0051] Replacer **89** is configured to replace a co-located portion of the renderable portion **46** of the reconstructed version **58** of the first view **24** as obtained by reconstructor **84**, by this actually rendered portion **57** so as to obtain a mixed signal **91**.

[0052] The mixed version **91** could represent the presentation output of the decoder of FIG. 5. However, alternatively, other portions beyond those rendered in the mixed signal, namely by signal **57**, could further be replaced by further actually rendered portions in renderer **89'**.

[0053] The reference signal **59** of view predictor could therefore be composed of the reconstructed signal **58**, the mixed signal **91** and the rendered reference signal **57**. Naturally, one of the mixed signal **91** and the rendered references signal may be missing.

[0054] In the following, embodiments are described as to how the encoder may inform the decoder as to which portion is to be actually rendered, which of the actually rendered portion is to be introduced into the prediction loop, either by replacement or in juxtaposition to the reconstructed signal **58**, and which of the actually rendered portion is to be introduced into the finally reconstructed output **92** to be reproduced, or how this information may otherwise derived by the decoder. For example, actually rendered portion may correspond, by default, to the whole renderable portion **46**, and be used as rendered reference signal **57**. Alternatively, the data stream generator **44** may be configured to encode a position information indicating a position of the renderable portion **46** into the data stream. Further. position information indicating portions of the multi-view signal for which the rendered reference **57** is—via mixed signal **91** or not—to be used, may be coded into the data stream in form of, for example, one or more special coding modes. IN this regard, reference is made to the below description. Summarizing the above embodiments for an encoder and a corresponding decoder, the data stream **21** according to the above embodiments may comprise: 1) residual data for the non-renderable portion **48** and, optionally, renderable portion **46**; 2) prediction parameters for both blocks within the renderable portion **46** as well as blocks within the non-renderable portion **48** or, optionally, merely within the non-renderable portion **48**; some of these prediction parameters would be motion and/or disparity compensated parameters dedicated for copying portions of the reference signal stemming from the rendered signal **57** (either directly, or via the mixed signal); 3) optionally, subdivision information for subdividing the views into the blocks, 4) optionally, coding modes associated with the blocks and determining the sort of prediction parameters associated with the respective block. Regarding the prediction parameters, same may at least partially predictively coded into the data stream in the form of prediction parameter residuals. At least some of the prediction parameters relating to blocks within the non-renderable portion may be coded into the data stream by way of prediction parameter residuals relative to prediction parameters relating to the renderable portion **46**. Naturally, some of the prediction parameters of the non-renderable portion may also be coded into the data stream as prediction parameter residuals relative to prediction parameters relating also to the non-renderable portion **48**. Similarly, some of the prediction parameters concerning the renderable portion **46** may be coded into the data stream in the from of a prediction parameter residuals relative to prediction parameters of other blocks within the renderable portion **46**.

[0055] After having described, rather generally, embodiments of the present invention, some possible variations are described below. For example, as has already been described above, the view predictor **38** could alternatively perform the block based prediction from a reference signal that, at least partially, is composed of the original version of the previously encoded portion of the multi-view signal **20**. However, the reconstructible version **58** is advantageous. Similarly, although FIG. 4 suggests that determinator **36** derives the position of the renderable portion **46** of view **24** based on the original version of view **22**, this is merely one option and it is advantageous that determinator **36** performs the determination based on the reconstructible version of view **22** as contained within reference signal **58**. This is illustrated by the dashed box **92** in FIG. 4 where the renderer **92** would also assume responsibility for the determination of portion **46**.

[0056] Thus, as become clear from FIG. 4, the reference signal **59**, from which view predictor **38**

performed the block-based prediction, i.e. the motion compensated and/or the disparity compensated copying, is not only defined by the reconstruction result **58** derived by constructor **56**, i.e. the reconstructed version of multi-view signal **20** as currently reconstructed. Rather, even versions rendered therefrom may contribute to the reference signal **59** based on which view predictor **38** performs the block-based prediction. In accordance with a first alternative, for example, the determinator **92** is configured to actually render the renderable portion **46** to obtain a rendered signal. The rendered signal **57** within the renderable portion **46** forms, along with the reconstructed signal **58** within the non-renderable portion **48**, a mixed reference signal **91** based on which the view predictor may perform the block-based prediction. That is, in a decoded picture buffer **94**, optionally positioned between renderer **92** and view predictor **38**, the reference pictures stored therein would correspond to the reconstructed version **58** as reconstructed by reconstructor **56** as far as the non-renderable portion **48** within these pictures is concerned, and the rendered signal **57** as rendered by renderer **92** as far as the renderable portion **46** within these pictures is concerned. That is, the motion- and/or disparity-compensated prediction performed by view predictor **38** would be performed on reference pictures corresponding to those reproduced at the decoding side. A decoded picture buffer **94'** may likewise be present in the prediction loop in the decoder right in front of the prediction loop input of the view predictor **82**, with same being filled in exactly the same manner as at the encoder in the way set out hereinafter for further embodiments.

[0057] In accordance with an even further alternative, the determinator **92** may be configured to actually render the renderable portion **46** to obtain a rendered signal **57** from the reconstructed version **58** of view **22**, and the decoded picture buffer **94** would buffer two kinds of reference signals, namely the reference signal **58** as obtained by reconstructor **56** having the non-renderable as well as renderable portions reconstructed using the prediction signal **52** and residual signal **54**, and a further reference signal in form of the rendered signal **57** within the renderable portion **46**. Thus, different reference indices for the motion—and/or disparity—compensation predicted blocks would refer to different kinds of pictures, namely pictures as output by reconstructor **46** on the one hand, and pictures being rendered at the renderable portion **46** on the other hand. Both renderers **89** and **89'** could be left away in that case.

[0058] In accordance with an even further alternative, the determinator **92** may be configured to actually render the renderable portion **46** to obtain a rendered signal **57** from the reconstructed version of view **22** for specific blocks for which the reconstruction based on rendering from view **22** is activated by a respective coding mode option whereas other blocks are reconstructed based on the reconstruction result of reconstructor **56**. The result would be the mixed signal **91**.

[0059] After having described different embodiments of the present invention, further possible variations and further details regarding the encoding procedure are set out below.

[0060] Further, according to the possible implementation of the embodiments outlined below, a more detailed description is spent on a possible way to signal within the data stream **21** the omission of the residual signal within the renderable portion **46** in accordance with some embodiments of the present invention. As will be outlined in more detail below, the prediction modes for the blocks may signal to the decoder that a residual signal for the respective block is missing or not. In that case, based on the coded depth maps of the reference view **22** and the coded camera parameters, which may be additionally coded as an optional element **5**) in data stream **21**, the encoder **26** (or determinator **92**) determines regions **46** in the currently coded image at time instance **T** that can be rendered from images of any previously coded reference view **22** of the same time instance **T**. In order to determine such regions **46**, the encoder/determinator can use basically the same procedure that is used in the renderer **86** for determining the regions that are rendered from a particular view.

[0061] For the regions **46** that can be rendered, no residual information may be coded. I.e., for these regions **46**, the coded block flags or coded block pattern, which indicates whether non-zero

transform coefficient levels are present, may be switched off, such as set to zero at the encoder side. However, splitting information for the prediction tree (which signal how a block is decomposed for the purpose of temporal or inter-view prediction) and motion data may still be transmitted for these regions, in order not to break the chain of motion data prediction. As described above, when the motion information would be skipped or zero motion vectors would be transmitted for these regions, a significant amount of bit rate would have to be spent for transmitting the motion data of the next block that lies in a region **48** that cannot be rendered. There would not be any suitable motion data in spatially or temporally neighboring blocks that can be used for efficiently predicting the motion data of the next block that cannot be rendered. By also transmitting the motion data of blocks that can be rendered (perhaps with less accuracy, as described below), this can be prevented and overall bit rate may be saved. Thus, e.g. in block-based hybrid video coding, blocks containing only samples that can be rendered are simply encoded without residual information.

[0062] If regions **46** that are identified to be renderable are coded without residual information, these regions can only be reconstructed in a lower quality. Thus, these regions should be rendered **86** from previously decoded views after decoding. In an embodiment, these regions are identified at the decoder side by using a particular specified algorithm (the same or a similar algorithm that is used in the encoder for identifying these regions) in a renderer **86**. In another embodiment of the invention, these regions are indicated in the bitstream, e.g. by transmitted supplemental enhancement information (SEI message).

[0063] Thus, as became clear from the just-described implementation possibility, it could be that the omission of a residual signal within data stream **21** is realized at block granularity, i.e. by transmitting coding modes for the blocks of view **24** such that for blocks completely residing within the renderable portion **46**, no residual signal information is present within the data stream **21**, with this fact being known for the decoder from the respective coding mode. With regard to the other blocks, i.e. blocks partially overlapping at the renderable portion and partially overlapping the non-renderable portion, the residual signal may be transmitted within data stream **21** completely, or merely with respect to that part of these blocks lying within the non-renderable portion **48**.

[0064] If there are only blocks of a fixed size, the encoding can be simply done as described above. For each block, it can be tested in determinator **36** or **92** whether it can be rendered, and if it can be rendered, no transmission of any residual takes place. However, the video coding algorithm used in FIGS. **4** and **5** may include a mechanism by which a (large) block can either be coded in a particular coding mode with associated motion parameters or it can be split into smaller subblocks and each of these subblocks is transmitted in a particular coding mode. In order to decide whether a block, which cannot be completely but only partly rendered, is split or not, encoders may have to compare any particular cost measure. For example, when using a recursive partitioning, such as quad-tree structures, the RD-costs of different tree-depth need to be compared.

[0065] Rendering artifacts have a different impact to the subjective image/video quality perception than coding artifacts. Therefore, distortions introduced by rendering cannot be compared directly to distortions introduced by lossy compression. However, in typical encoders, the coding mode of a block is usually determined by minimizing a distortion measure D or a Lagrangian cost measure $D + \lambda \cdot R$, which weights the distortion D for a coding mode against the associated number of bits R given a Lagrangian multiplier λ . If the mode decision is based on such a cost measure in encoder **26**, only a very low percentage of the blocks would be coded in the inter-view skip mode described above, i.e. the mode with suppression of residual coding, since the distortion would usually be quite high compared to conventional coding modes, although the final rendered result will look fine.

[0066] Due to that reason, it is advantageous to modify the distortion calculation in the mode decision process. In an embodiment of the invention, the distortion is calculated by only considering samples that cannot be rendered. For all samples that can be rendered (as determined

by the determinant), the sample distortion is that to zero (independent of the coding mode) so that it doesn't have any impact on the overall distortion for a block. Therefore, the distortion of rendered regions is not considered in the mode-decision.

[0067] FIG. 6 shows an example: The right image shows the block subdivision that is one level deeper than in the left image. The shaded area labels the samples that can be rendered and are therefore not taken into account. Thus, for example, the upper left block D.sub.1 in the right image is not considered at all. And when comparing the cost for the entire block D.sub.0 with the sum of the costs for the remaining subblocks D.sub.1+D.sub.2+D.sub.3, none of the shaded samples (which can be rendered) is taken into account. By this modification, we do not automatically split blocks for which a subblock can be rendered, but may also code the entire block D.sub.0 using a conventional coding mode if this improves the overall coding efficiency.

[0068] Thus, the residual generator **40** and the view predictor **38** could be configured to locally vary a block granularity at which the first and second prediction parameters are defined, and at which the residual signal **54** covers the dependent view **24**. In that case, both could cooperate to signal the local variation of the block granularity within the data stream **21** in form of the above-mentioned subdivision information. The residual generator and the view predictor could then be configured to perform the local variation at a boundary of the renderable view, such as the block-subdivision of the block **Do**, using a minimization algorithm so as to minimize a cost measure which increases depending on a distortion measured outside the renderable portion **46**, i.e. within portion **48**, stronger than on a distortion measured inside the renderable portion **46**.

[0069] However, even a scaling of λ in the Lagrangian cost measure, i.e. the Lagrangian multiplier λ , could be scaled in the calculation of the RD cost measure within the renderable portion compared to the non-renderable portion.

[0070] As described above, for blocks which can be rendered, no transmission any residual information takes place, but splitting and motion information is still transmitted in order to improve the overall coding efficiency for the motion parameters (see above). These motion parameters are, however, only necessitated for obtaining a suitable prediction for the motion parameters of following blocks, which cannot be rendered completely. It is usually justified that these splitting and motion parameter are coded with less accuracy. I.e., when selecting the splitting and motion parameters, the associated rate R should have a larger influence on the selected parameters than for conventionally coded blocks. Or in other words, since the content of the regions **46** that can be rendered is already present in the bitstream (in a different view), these regions do not have to be encoded in a high quality and the calculation of R-D costs may be adapted towards a weaker influence of the distortion term. This forces the mode decision to modes of lower bit rate. Due to that reason, the calculation of the RD-costs is changed from $C=D+\lambda R$ (with R being the rate, D being the distortion and λ being the Lagrange Multiplier) to $C=D+s \lambda R$. The Lagrange Multiplier λ is scaled by a factors $s>1$.

[0071] That is, the encoder **26** may be configured to set prediction parameters relating to the renderable portion **46** using a minimization algorithm so as to minimize a first cost measure, and set prediction parameters relating to the non-renderable portion **48** using the minimization algorithm so as to minimize a second cost measure, wherein a quotient of a strength of an increase of the first cost measure depending on bit rate, divided by a strength of an increase of the first cost measure depending on distortion is greater than for the second cost measure. In the example presented above, the quotient would be **SA** within the renderable portion, and **2** within the non-renderable portion.

[0072] Focus is now put onto the above emphasized possibility to include the actually rendered signal into the prediction loop so as to form the prediction basis for view predictor **38**. The decision whether the renderable portion **46** is included into the prediction loop or not, could be performed on a block basis as described further below. In particular, in accordance with this embodiment, the inter-view skip mode could be realized by signaling the rendered region and integrating the

rendering into the prediction loop via renderer **92**. The bitstream syntax and the decoding process are modified. In contrast to above embodiments, it is signaled inside the bitstream **21** whether a block is rendered based on data coded in an already coded view **22** and, for these blocks, the reference signal **59** entering view predictor **38** is obtained by actually rendering these blocks inside the decoder loop within renderer **92**. Consequently, the rendered blocks are also used when these blocks are referenced by future frames for the purpose of motion compensation. That is, their content is rendered when copied to another block.

[0073] Thus, based on the coded depth maps of the reference view **22** and the coded camera parameters, the encoder/determinator determines regions **46** in the currently coded image of view **24** at time instance T that can be rendered from images of previously coded reference view **22** of the same time instance T. This process is the same as described above. But additionally, for particular blocks in the prediction tree of the currently coded image of view **24** a flag is transmitted inside the bitstream, indicating whether the block can be rendered or not. For those blocks that can be rendered, no residual information is transmitted. The flag could be transmitted for all blocks that are coded using motion (or disparity) compensation. In another embodiment, the flag is only transmitted for a subset of the blocks (for example only blocks with a block size larger than or equal to a particular block size); for all other blocks, the flag may be inferred to be equal to 0 (i.e., these blocks are coded using conventional coding mode).

[0074] The reconstruction of these partitions is obtained by the particular rendering algorithm in renderer **92**. If such a rendered block is used for motion-compensated prediction of other blocks in following frames of view **24**, for example, the rendered reconstruction signal is used.

[0075] In an embodiment, the described possibility, in which a flag signals whether a block is rendered and the reconstruction signal is obtained by actually rendering the blocks, is combined with the modified encoder control described above.

[0076] However, the explicit signaling whether a certain block within the renderable portion **46** is to be replaced within the prediction loop by the rendered signal does not need to be transmitted explicitly. Rather, as described above this replacement may be contacted automatically based on the determined location of the renderable portion **46**. Thus, the inter-view skip mode may be realized by integrating rendering into the prediction loop, but without transmitting block-based rendering information. In that case, the bitstream syntax does not include a block-based indication whether a respective block is rendered or not, but the decoding process is modified in a way that the reconstruction signal of blocks that can be rendered is obtained by actual rendering. This method is in principle similar to the just-described possibility of signaling block-wise the insertion of the rendered version or reconstructed version into the prediction loop, but no flags are transmitted that indicate if the region can be rendered. Instead, the decoder checks renderability itself using the received depth maps and camera parameters the same way as the encoder does, and for renderable blocks the use of the rendered version from the reference version is used by default.

[0077] Based on the coded depth maps of the reference view **22** and the coded camera parameters, the encoder/determiner determines regions **46** in the currently coded image of view **24** that can be rendered from images of previously coded reference view **22** of the same time instance. For those regions **46**, no further residual information is transmitted. The reconstruction of these regions is obtained by rendering. In an embodiment of the invention, this method is combined with a modified encoder control as described above.

[0078] Based on the decoded depth maps of the reference views and the decoded camera parameters, the decoder/renderer **86** determines regions in the currently decoded image that can be rendered from images of previously decoded reference views of the same time instance. The reconstruction of these regions is obtained by rendering. The algorithm that is used for determining the regions that can be rendered is the same in encoder in decoder, so that no encoder-decoder mismatches can occur if the bitstream is transmitted without errors.

[0079] As already described above, a kind of warped inter-view prediction may be realized. For

reducing the bit rate associated with blocks in view **24** that can be rendered based on the pictures and depth of an already coded reference view **22**, new rendered reference pictures may be introduced and buffered in buffer **94**. An image is created by rendering the current view **24** using the reconstructed image of one or more reference views **22** and the associated depth maps **28**. For that purpose, the images of an already coded view **22** are mapped to the coordinate system of the current view by using the associated depth maps by renderer **22**. The newly created image represents a suitable representation for the current image of the current view **24**, except for the areas **48** that are uncovered in the current view **24**, or are not present in the reference views such as the border area illustrated to the right of texture **32**. During the rendering process, the holes in the rendered image can be either set to a particular values or can be filled by particular algorithms, e.g., by setting the samples equal to the sample values of the neighboring samples with the smallest depth value, a procedure called background filling.

[0080] In an embodiment, a single rendered image may be created based on the data of all previously coded views (for the current time instance) or a subset of the previously coded views. In another embodiment of the invention, a rendered image is created by using exactly one previously coded view. However, if more than one view has already been coded, two or more rendered images can be created. In another embodiment, a rendered image can be created by using the data of exactly one view or by using the data of more than one view. It can be signaled inside the data stream, which data are used for creating a particular rendered image.

[0081] In an embodiment, the rendered images may replace the interview reference picture for the same reference view in the reference picture lists of the current picture. In another embodiment, both conventional interview reference pictures and rendered pictures that are rendered using the reconstructed pictures and depths maps of already coded views can be presents in the reference picture lists or lists for the current picture. What reference pictures are actually included in the reference picture list(s) may be signaled inside the data stream.

[0082] In an embodiment, the low-level syntax is not modified, i.e., only the generation of the reference picture list is changed (including the creation of rendered reference pictures). The advantage of this method is that, when an interview reference picture is used (as signaled by the reference picture index), the associated motion parameters are zero or at least nearly zero (since the reference view image has been mapped to the coordinates of the current view) and the bit rate associated with the motion data is very small. In order to further reduce the bit rate associated with regions that can be rendered, the encoder control can be modified in a similar way as described above. That means, for block that can be rendered (as determined by the encoder), the encoder can set the reference index to a value that indicates the usage of a rendered reference picture. In addition, the residual information can be set to zero (i.e., no residual is coded for a rendered block). Also, the motion vector can be set to zero. Furthermore, the cost measure for deciding the coding mode for a block that can be rendered can be modified by multiplying the Lagrangian parameter with a scaling factor greater than 1. Also, during mode decision, the distortion measure can be modified in a way that only samples that cannot be rendered contribute to the distortion.

[0083] In another embodiment, the motion vector differences, i.e. the prediction parameter residual, are not transmitted but inferred to be equal to 0 if the reference index indicates that a rendered reference image is used (but the residual can be transmitted). In a further embodiment, no residual and no coded block indication is transmitted if the reference index indicates that a rendered reference image is used (but motion vector differences can be transmitted). Instead, the coded block information is inferred to be zero (i.e., the corresponding residual is inferred to be zero). In another embodiment, no motion vector differences and no residual are transmitted if the reference index indicates that a rendered reference image is used.

[0084] In another embodiment, the concept of rendered pictures is not (or not only) applied to the coding of texture information (i.e., normal video signals), but it is applied to the coding of depth maps. In this case, the disocclusion areas of the rendered reference images can be filled with the

background depth (i.e., the depth value that specifies the farthest distance to the camera of the samples at the border of the disocclusion area). This hole filling can be applied line-wise (by considering only samples in the same line as the hole sample to be filled) or region-wise (by considering all samples at the borders of the disocclusion area).

[0085] After having described many different implementation possibilities, it should be noted that it became clear from the above discussion that renderer **86** within the decoder may not completely render all the renderable portion **46**. Further, block **86** may act as a determinator for determining the renderable portion with rendering portions thereof signaled within the data stream **21**.

[0086] Although some aspects have been described in the context of an apparatus, it is clear that these aspects also represent a description of the corresponding method, where a block or device corresponds to a method step or a feature of a method step. Analogously, aspects described in the context of a method step also represent a description of a corresponding block or item or feature of a corresponding apparatus. Some or all of the method steps may be executed by (or using) a hardware apparatus, like for example, a microprocessor, a programmable computer or an electronic circuit. In some embodiments, some one or more of the most important method steps may be executed by such an apparatus.

[0087] Depending on certain implementation requirements, embodiments of the invention can be implemented in hardware or in software. The implementation can be performed using a digital storage medium, for example a floppy disk, a DVD, a Blu-Ray, a CD, a ROM, a PROM, an EPROM, an EEPROM or a FLASH memory, having electronically readable control signals stored thereon, which cooperate (or are capable of cooperating) with a programmable computer system such that the respective method is performed. Therefore, the digital storage medium may be computer readable.

[0088] Some embodiments according to the invention comprise a data carrier having electronically readable control signals, which are capable of cooperating with a programmable computer system, such that one of the methods described herein is performed.

[0089] Generally, embodiments of the present invention can be implemented as a computer program product with a program code, the program code being operative for performing one of the methods when the computer program product runs on a computer. The program code may for example be stored on a machine readable carrier.

[0090] Other embodiments comprise the computer program for performing one of the methods described herein, stored on a machine readable carrier.

[0091] In other words, an embodiment of the inventive method is, therefore, a computer program having a program code for performing one of the methods described herein, when the computer program runs on a computer.

[0092] A further embodiment of the inventive methods is, therefore, a data carrier (or a digital storage medium, or a computer-readable medium) comprising, recorded thereon, the computer program for performing one of the methods described herein. The data carrier, the digital storage medium or the recorded medium are typically tangible and/or non-transitory.

[0093] A further embodiment of the inventive method is, therefore, a data stream or a sequence of signals representing the computer program for performing one of the methods described herein. The data stream or the sequence of signals may for example be configured to be transferred via a data communication connection, for example via the Internet.

[0094] A further embodiment comprises a processing means, for example a computer, or a programmable logic device, configured to or adapted to perform one of the methods described herein.

[0095] A further embodiment comprises a computer having installed thereon the computer program for performing one of the methods described herein.

[0096] A further embodiment according to the invention comprises an apparatus or a system configured to transfer (for example, electronically or optically) a computer program for performing

one of the methods described herein to a receiver. The receiver may, for example, be a computer, a mobile device, a memory device or the like. The apparatus or system may, for example, comprise a file server for transferring the computer program to the receiver.

[0097] In some embodiments, a programmable logic device (for example a field programmable gate array) may be used to perform some or all of the functionalities of the methods described herein. In some embodiments, a field programmable gate array may cooperate with a microprocessor in order to perform one of the methods described herein. Generally, the methods are performed by any hardware apparatus.

[0098] While this invention has been described in terms of several advantageous embodiments, there are alterations, permutations, and equivalents which fall within the scope of this invention. It should also be noted that there are many alternative ways of implementing the methods and compositions of the present invention. It is therefore intended that the following appended claims be interpreted as including all such alterations, permutations, and equivalents as fall within the true spirit and scope of the present invention.

Claims

1. A decoder for decoding a multi-view signal representing a video, comprising: a renderer configured for rendering, using a processor, a renderable portion of a first view of the multi-view signal, by view synthesis, based on a second view of the multi-view signal to generate a rendered portion of the first view; an extractor configured for extracting, from the data stream, using the processor, a first residual signal and a second residual signal; a view predictor configured for performing, using the processor, a block-based motion prediction based on the rendered portion and a non-renderable portion of the first view, at least by obtaining a first set of motion prediction parameters associated with the rendered portion, and generating a prediction signal based on the rendered portion, the first set of motion prediction parameters, the non-renderable portion, and a second set of motion prediction parameters associated with the non-renderable portion; a parameter predictor configured for obtaining, using the processor, the second set of motion prediction parameters based on the first set of motion prediction parameters and the second residual; and a reconstructor configured for reconstructing the first view based on the prediction signal and the first residual.
