



US 20250264858A1

(19) **United States**

(12) **Patent Application Publication**

DEAN et al.

(10) **Pub. No.: US 2025/0264858 A1**

(43) **Pub. Date:** **Aug. 21, 2025**

(54) **ASSET GROUP MANAGEMENT USER INTERFACE**

(71) Applicant: **Splunk Inc.**, San Francisco, CA (US)

(72) Inventors: **Erick Anthony DEAN**, Piedmont, CA (US); **Jindrich DINGA**, Los Altos, CA (US); **Marvin Herville GREEN**, Benicia, CA (US)

(21) Appl. No.: **19/201,315**

(22) Filed: **May 7, 2025**

Related U.S. Application Data

(63) Continuation of application No. 18/206,286, filed on Jun. 6, 2023, now Pat. No. 12,298,734, which is a continuation of application No. 17/374,239, filed on Jul. 13, 2021, now Pat. No. 11,703,826, which is a continuation of application No. 16/743,549, filed on Jan. 15, 2020, now Pat. No. 11,086,289, which is a continuation of application No. 16/147,439, filed on Sep. 28, 2018, now Pat. No. 10,564,622, which is a continuation-in-part of application No. 15/224,637, filed on Jul. 31, 2016, now Pat. No. 10,503,784, which is a continuation-in-part of application No. 15/224,641, filed on Jul. 31, 2016, now abandoned.

Publication Classification

(51) **Int. Cl.**

G05B 19/4063 (2006.01)

G06F 16/26 (2019.01)

G06F 16/901 (2019.01)

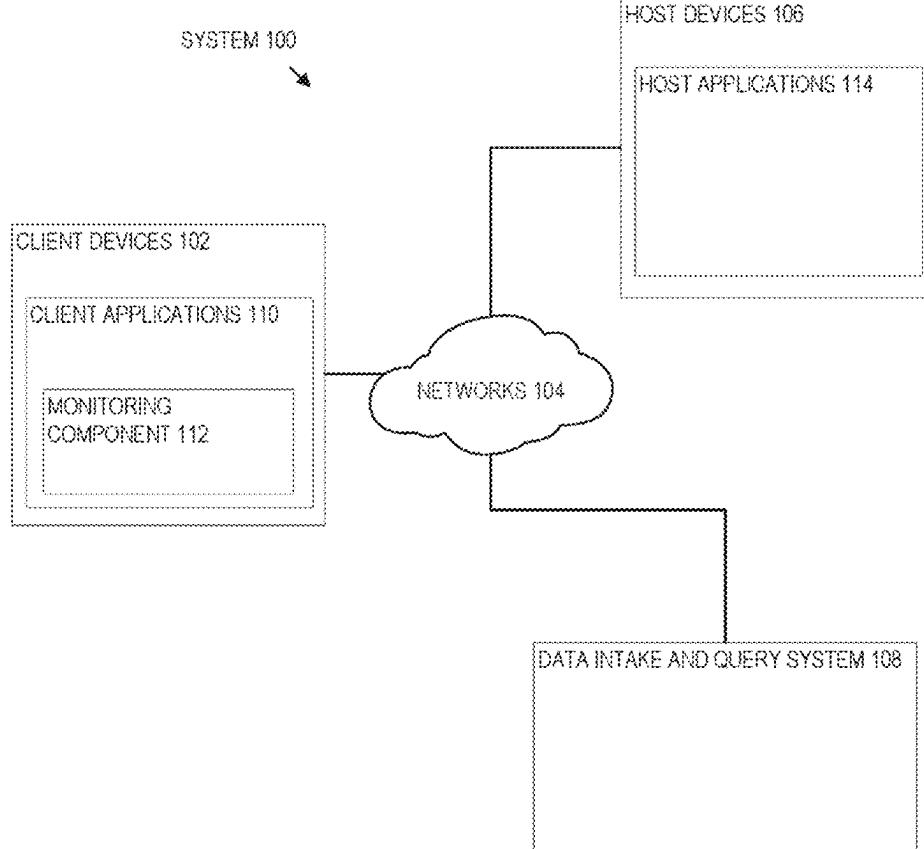
G06F 16/9535 (2019.01)

(52) **U.S. Cl.**

CPC **G05B 19/4063** (2013.01); **G06F 16/26** (2019.01); **G06F 16/9017** (2019.01); **G06F 16/9027** (2019.01); **G06F 16/9535** (2019.01); **G05B 2219/32404** (2013.01)

(57) **ABSTRACT**

A method of monitoring hierarchical asset groups includes the operations of: receiving a plurality of asset group configuration parameters comprising one or more association specifications, each association specification defining an association between an asset and a corresponding asset group of a plurality of asset groups of an asset hierarchy; receiving definitions of one or more asset group metrics, wherein each asset group metric of the one or more asset group metrics is derived from a search of machine data associated with a corresponding asset group; associating, with each asset group of the plurality of asset groups, a corresponding asset group metric of the one or more asset group metrics; generating one or more alert definitions associated with an asset group of the plurality of asset groups; monitoring, using the one or more asset group metrics, the asset hierarchy; and generating an alert based on the one or more alert definitions.



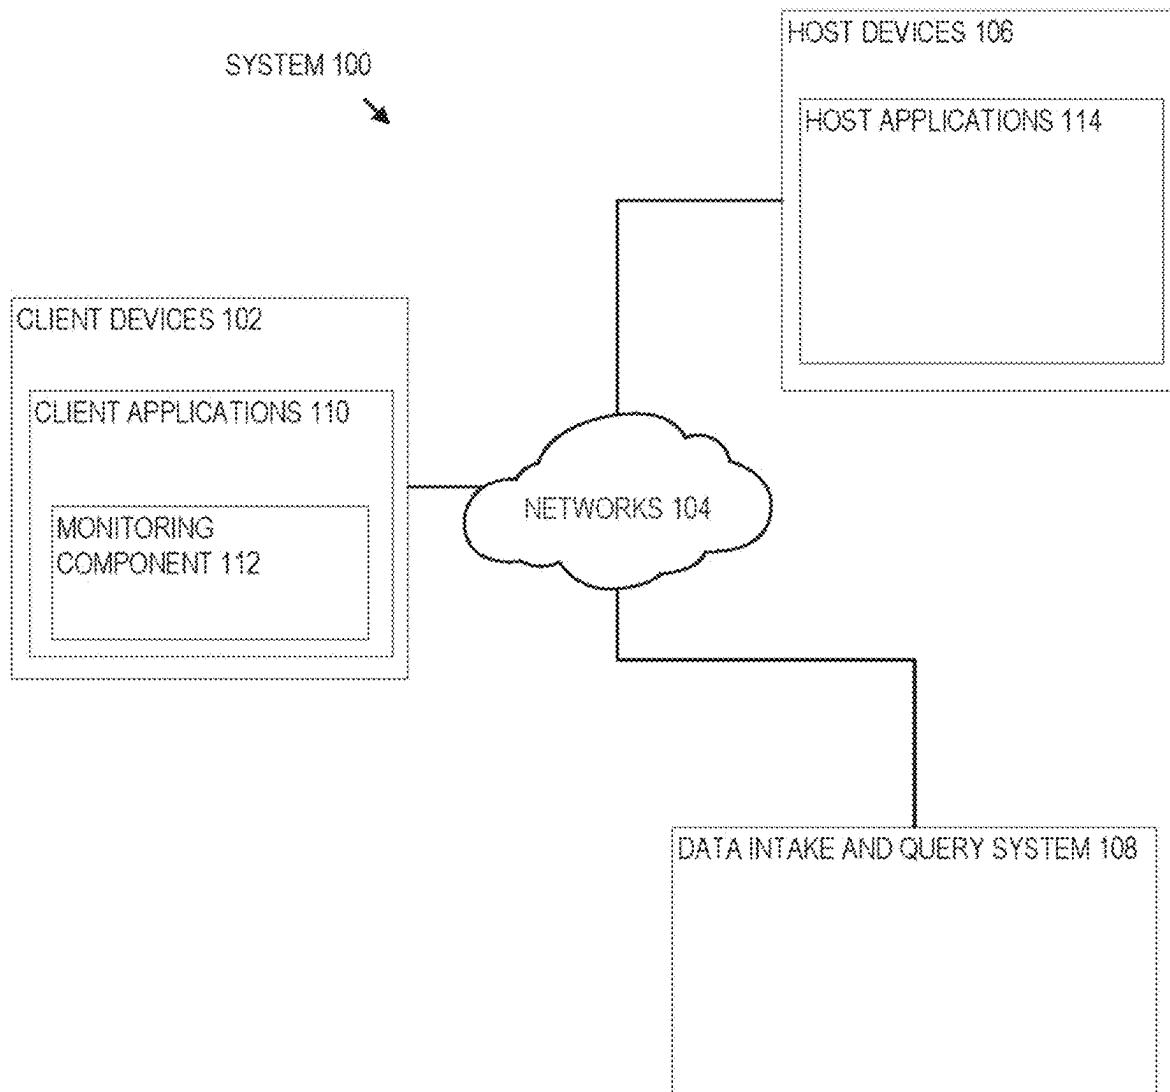


FIG. 1

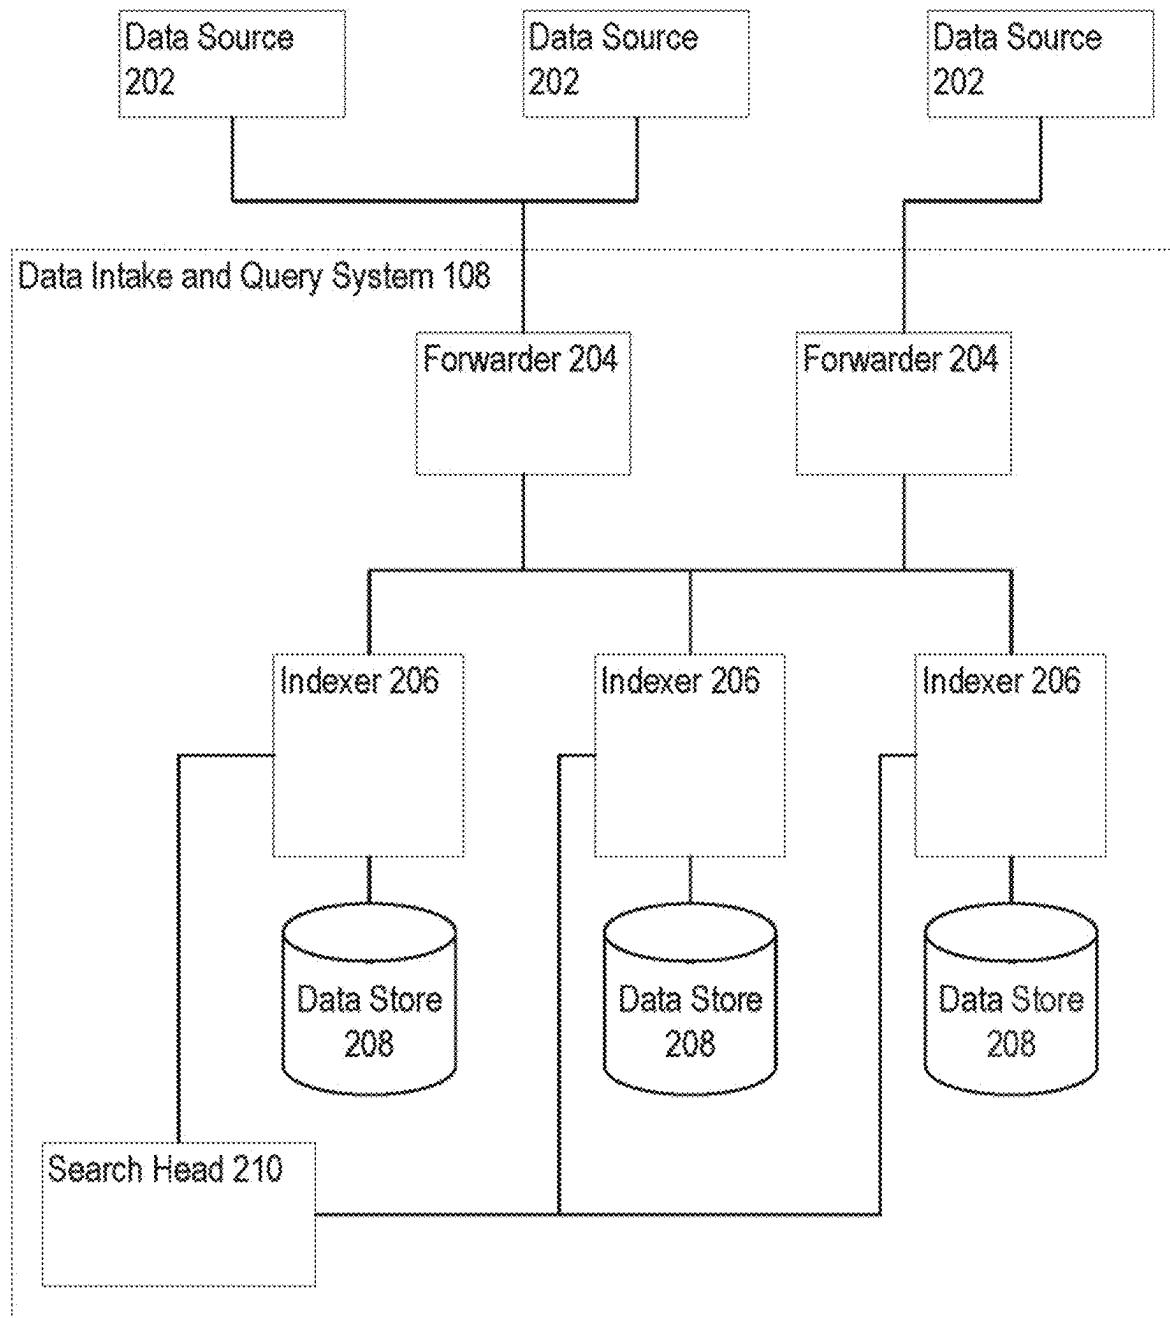


FIG. 2

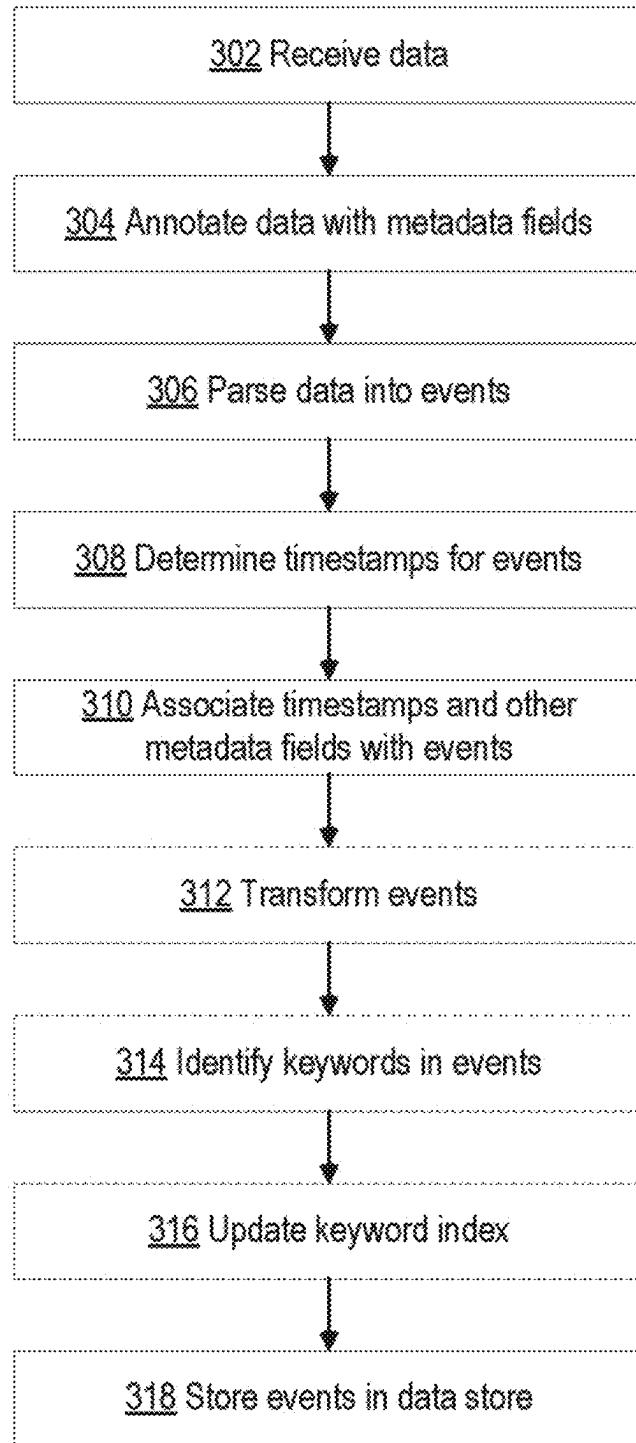


FIG. 3

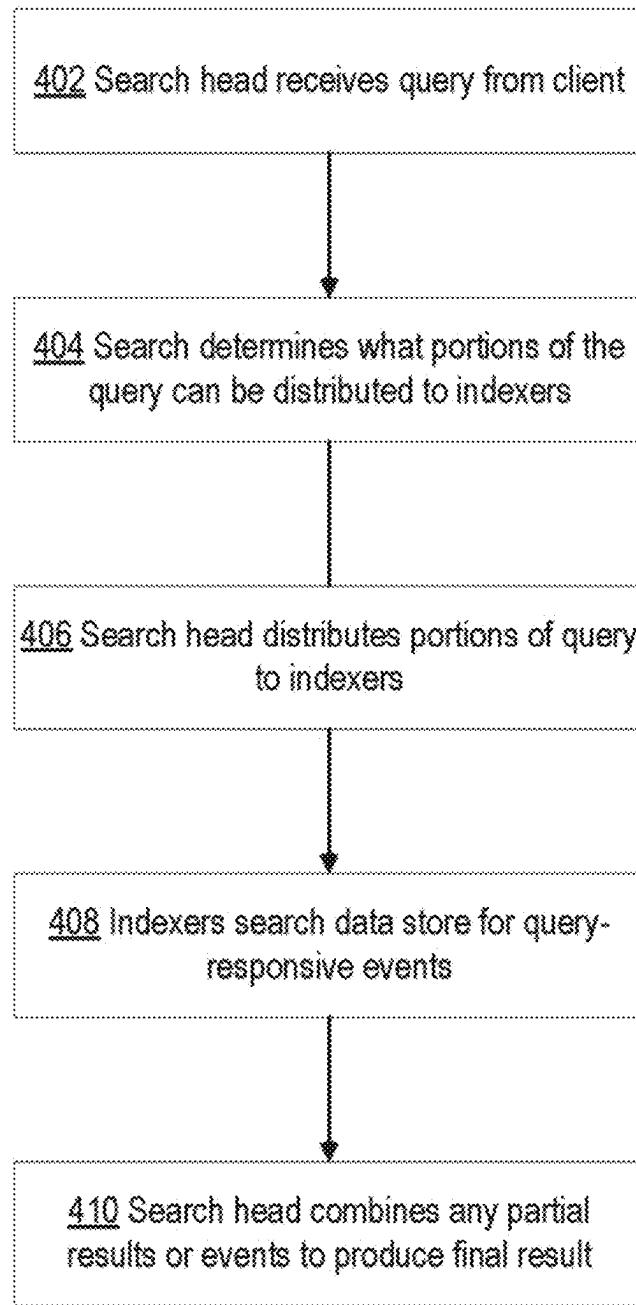


FIG. 4

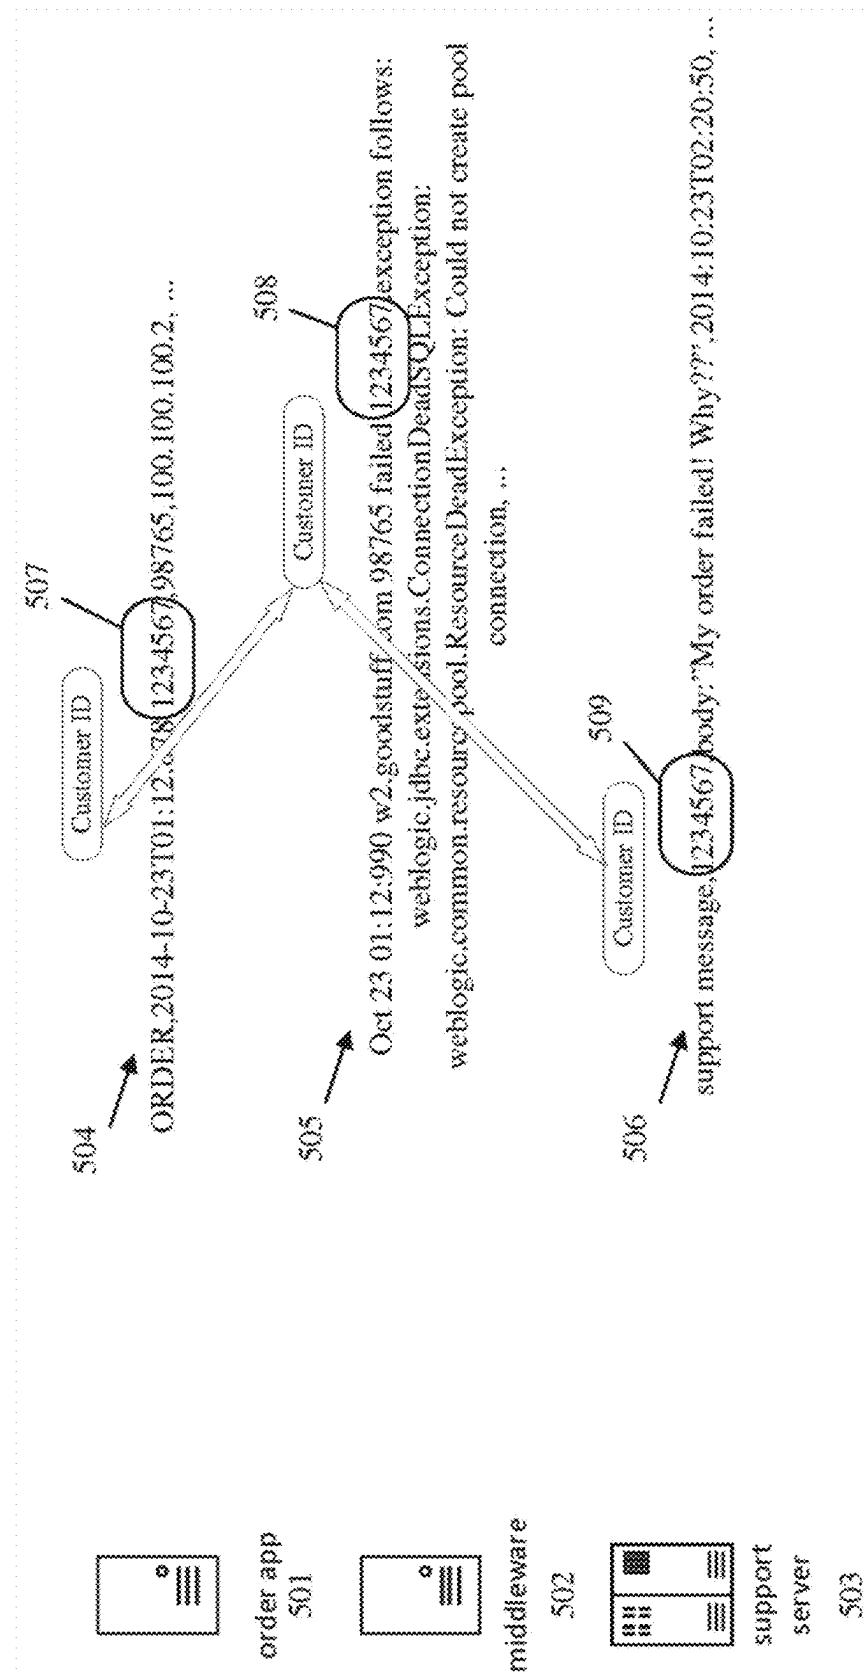


FIG. 5

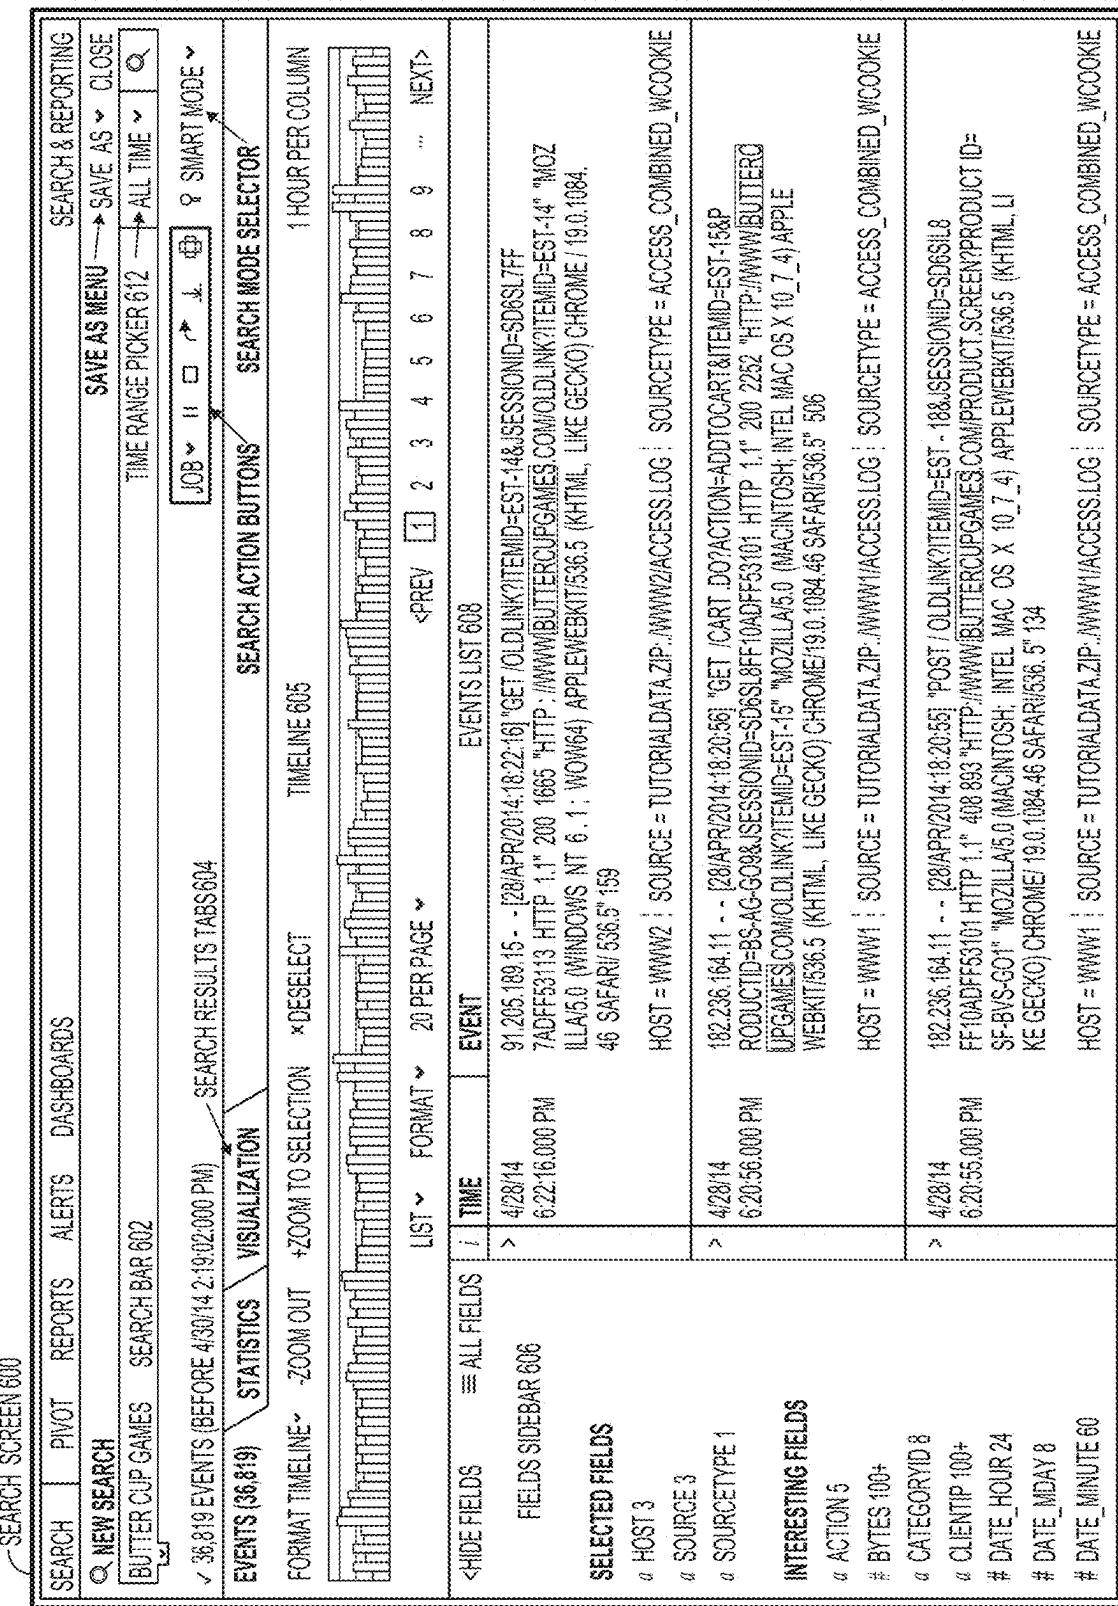


FIG. 6A

| HOSTS (5) \ SOURCES (8) \ SOURCETYPES (3) | | | |
|---|-----|---------|------------------------|
| FILTER | | | |
| HOST ◊ | ◊ | COUNT ♦ | LAST UPDATE ♦ |
| MAILSV | ◊ v | 9,829 | 4/29/14 1:32:47.000 PM |
| VENDOR_SALES | ◊ v | 30,244 | 4/29/14 1:32:46.000 PM |
| WWW1 | ◊ v | 24,221 | 4/29/14 1:32:44.000 PM |
| WWW2 | ◊ v | 22,595 | 4/29/14 1:32:47.000 PM |
| WWW3 | ◊ v | 22,975 | 4/29/14 1:32:45.000 PM |

FIG. 6B

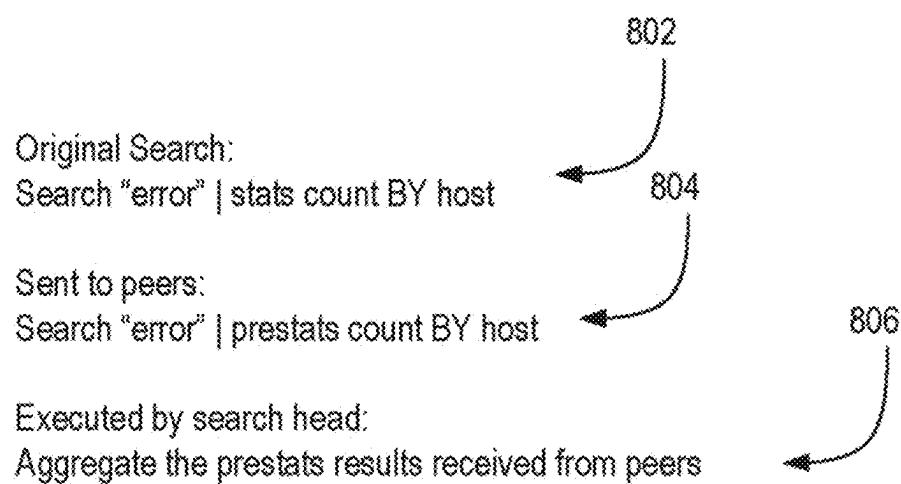


FIG. 7

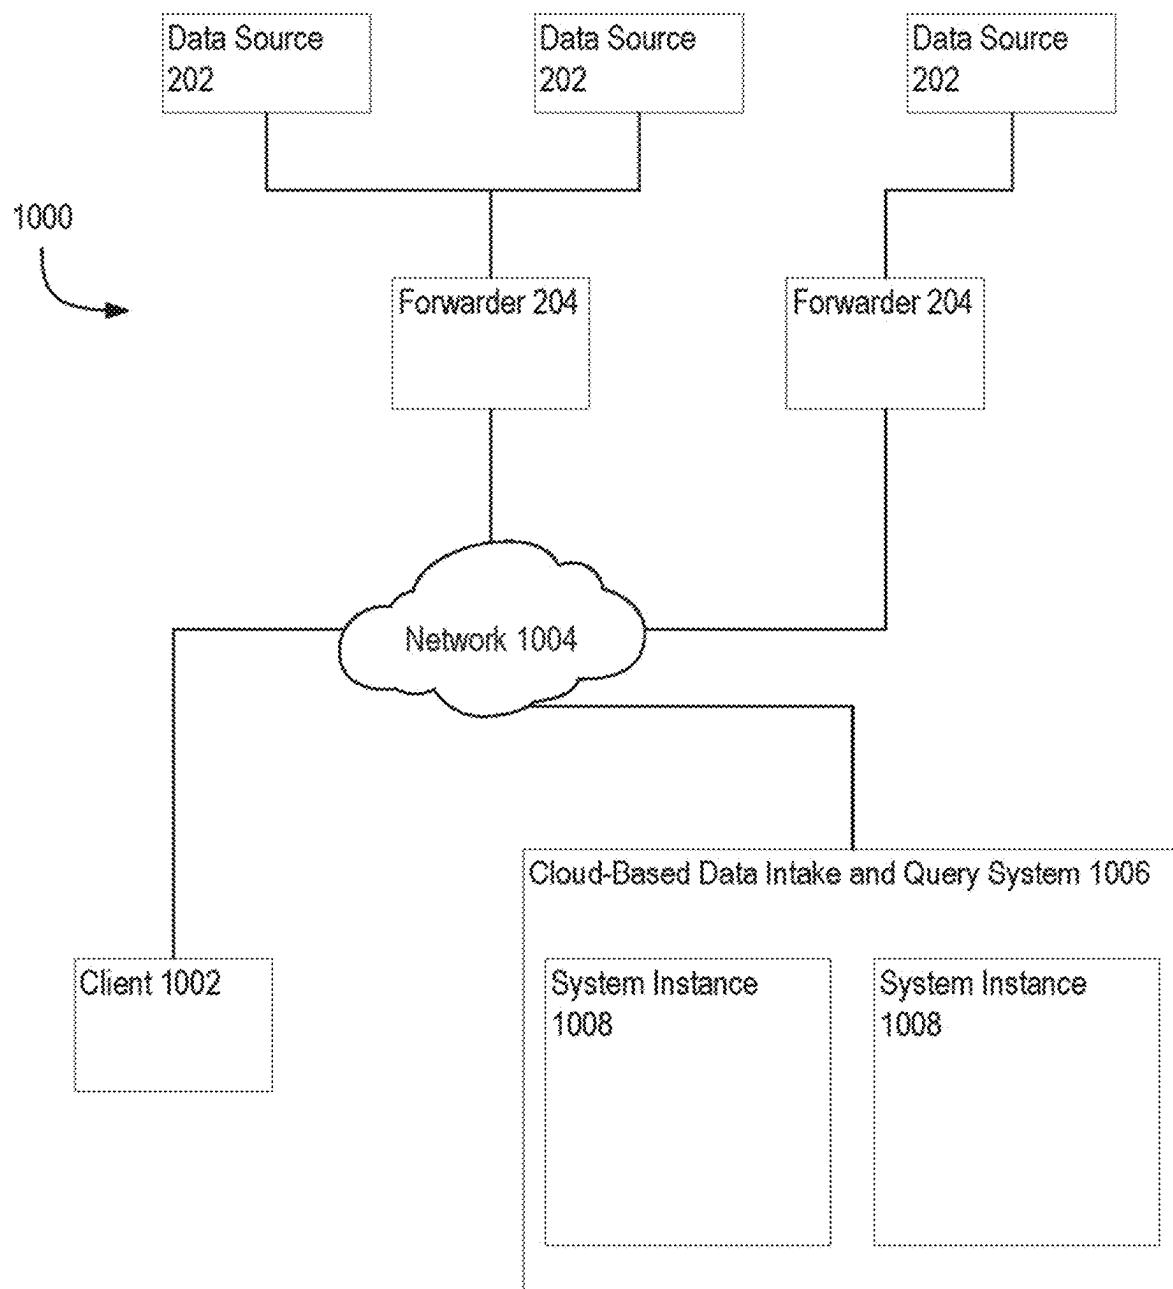


FIG. 8

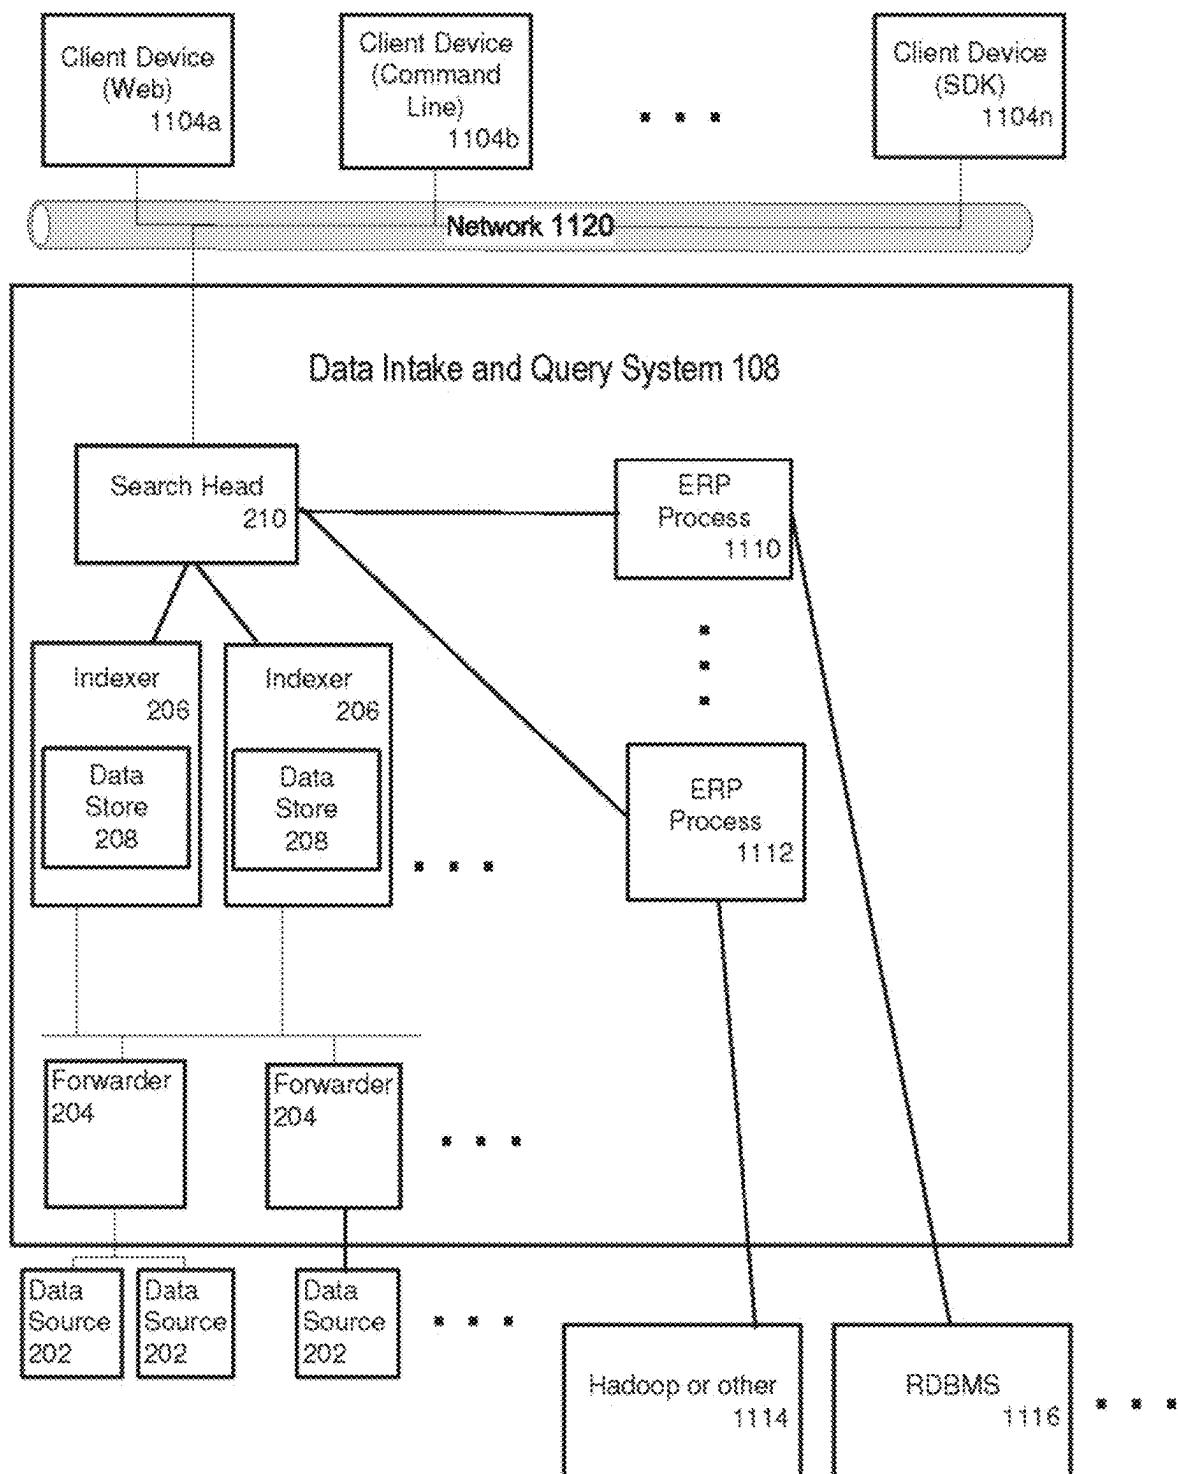


FIG. 9

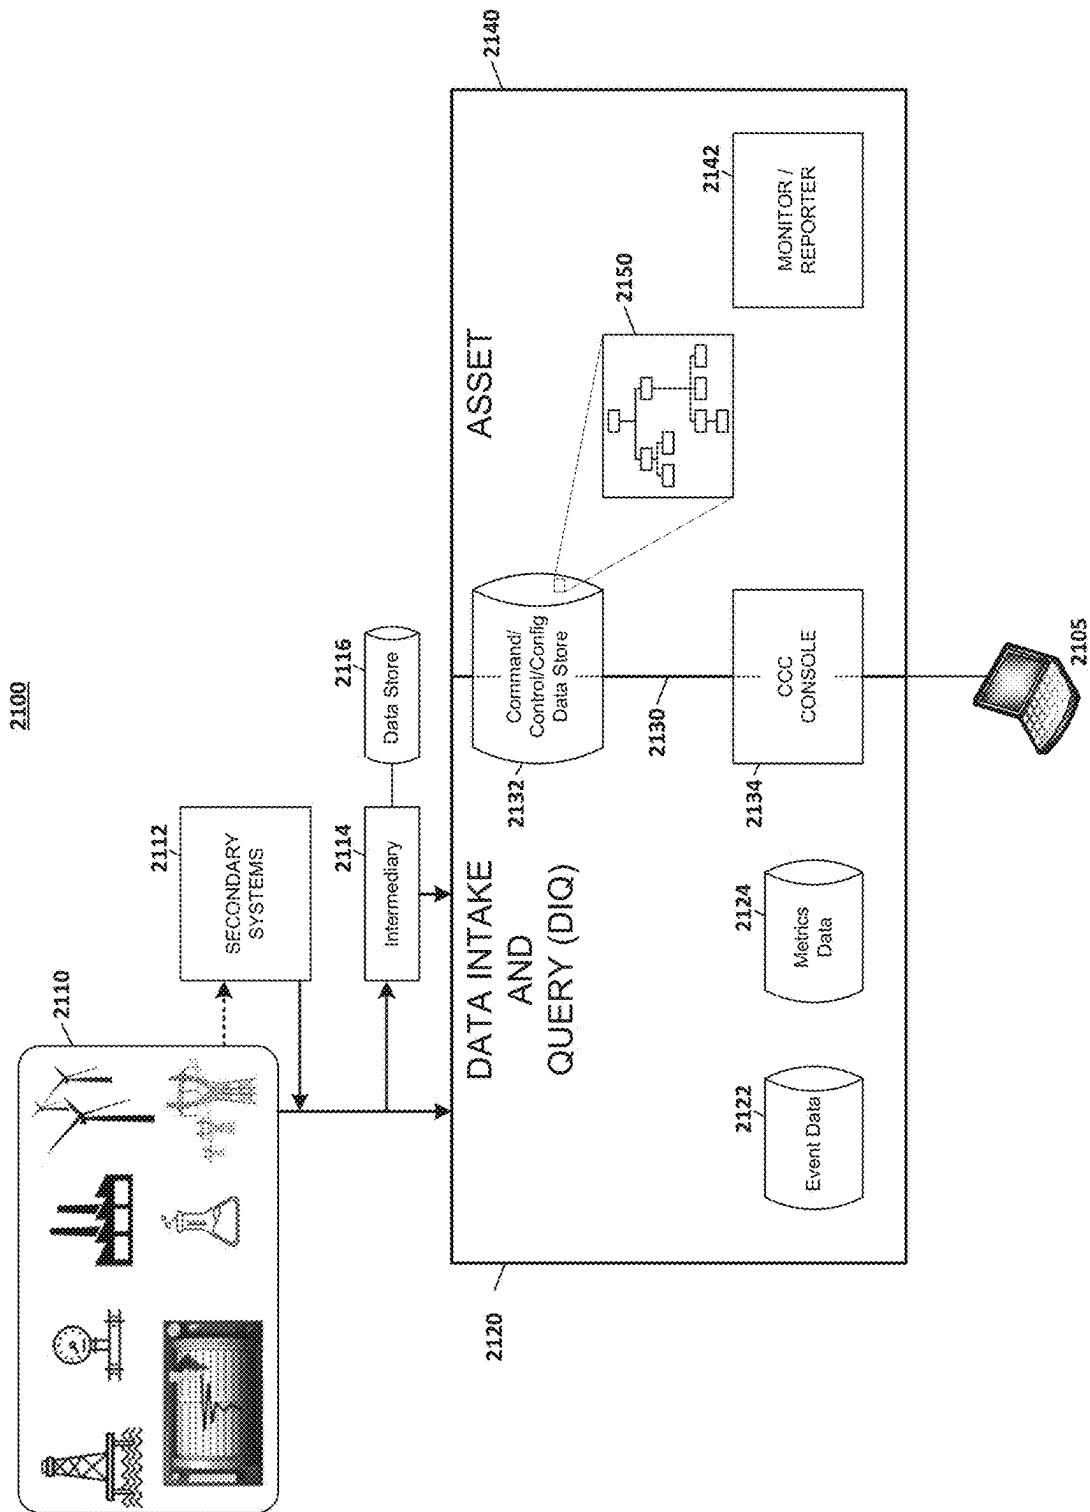


FIG. 10

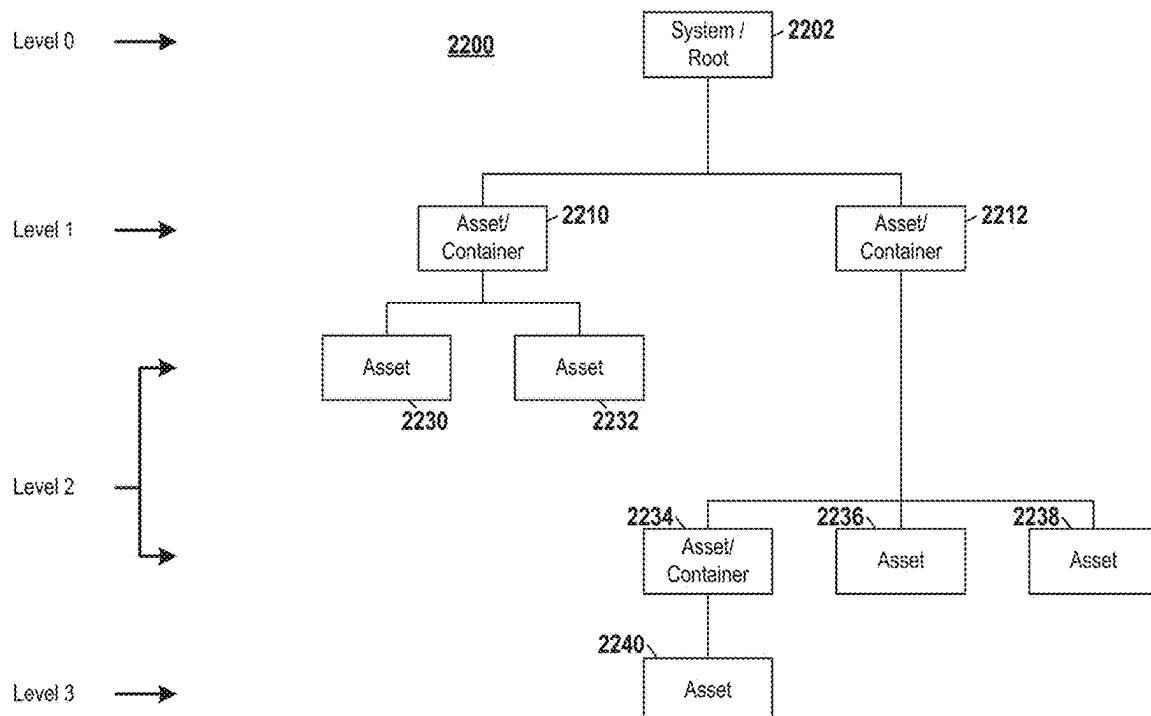


FIG. 11A

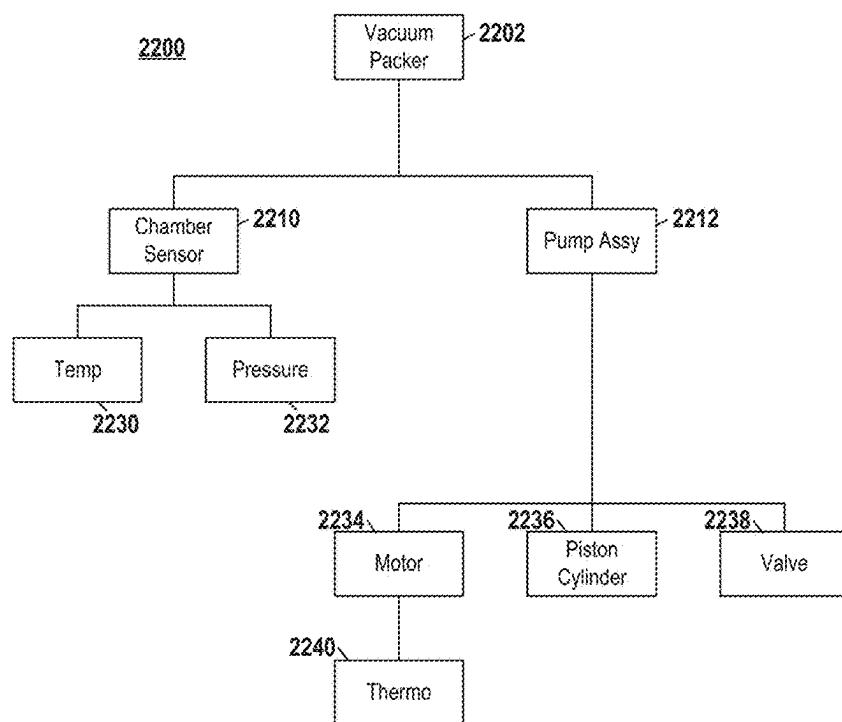


FIG. 11B

2300

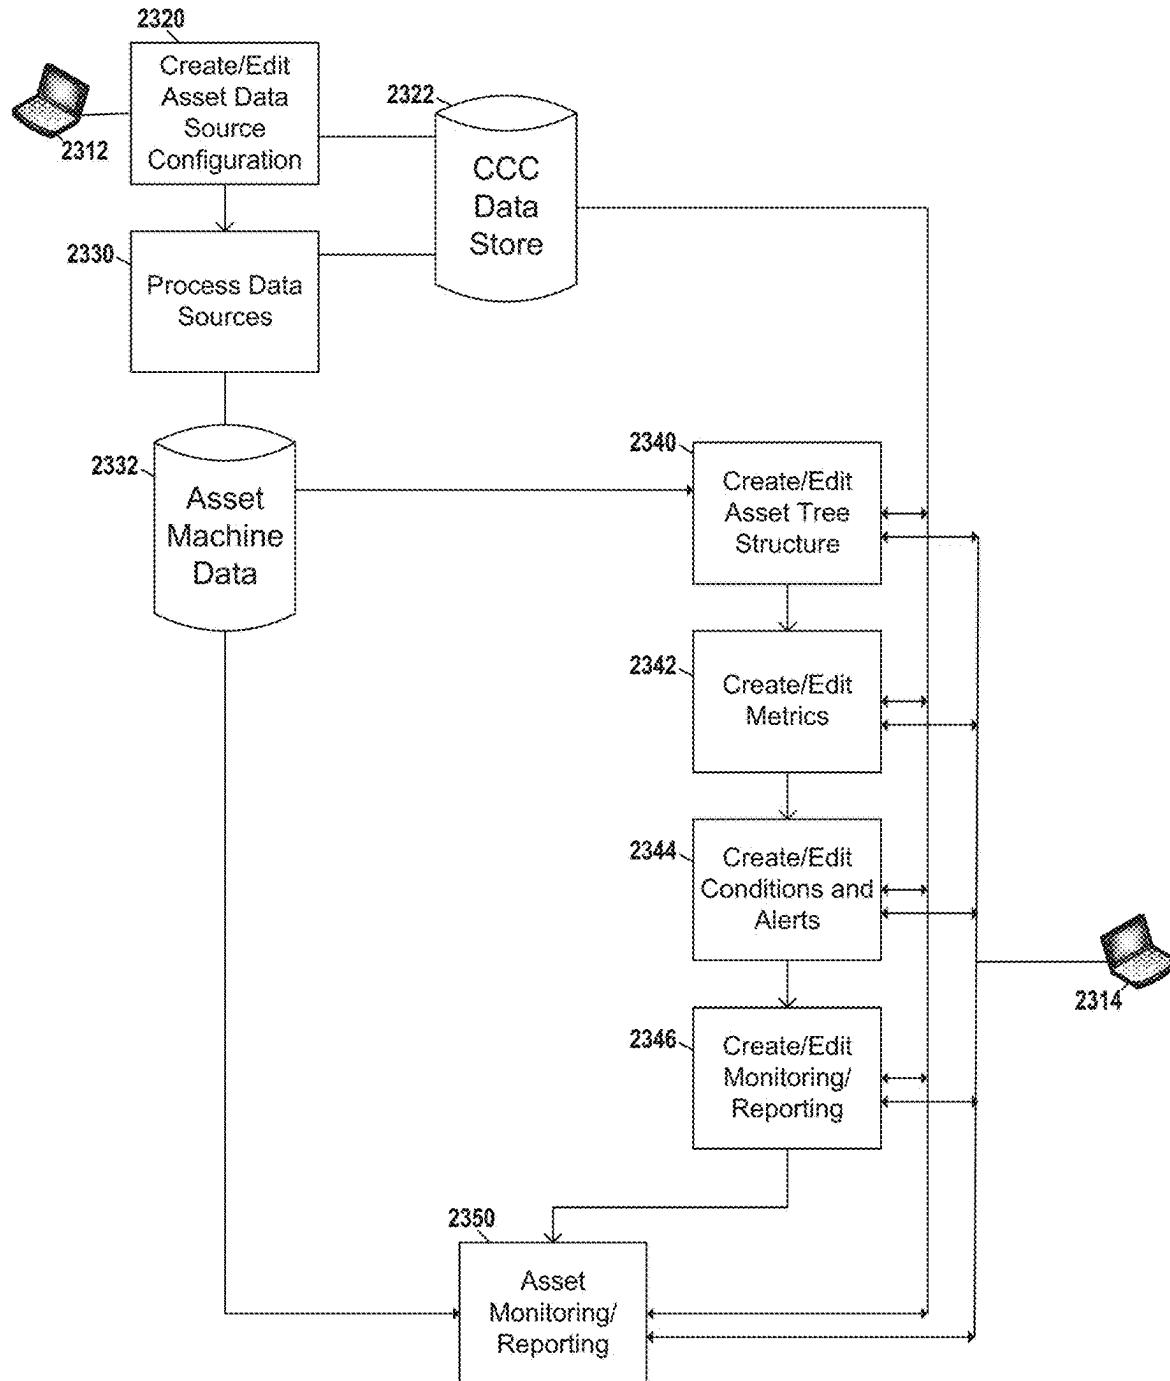


FIG. 12

| 2402 | 2404 | 2406 | 2408 |
|---------------------|---|---------|---------|
| 2410 | Data inputs | | |
| Local Inputs | | | |
| 2412 | See up date reports from files and databases, network ports, and scripted inputs. If you want to set up forwarding and receiving between two different instances, go to Forwarding and receiving. | Type | Type |
| 2418 | | Actions | Actions |
| 2420 | File & Directories Read a local file or monitor an entire directory. | Add New | Add New |
| 2422 | HTTP Event Collector Receive data over HTTP/Secure. | | |
| 2424 | TCP Generate TCP port for incoming data, e.g. reading. | | |
| 2426 | UDP Listen on a UDP port for incoming data, e.g. reading. | | |
| 2428 | Scripts Run execution scripts to collect or generate more data. | | |
| 2430 | OPC UA Pull Client/Read Manage the data subscription from OPC UA Server. | | |
| 2432 | OPC UA Event Notification Manage all event notifications from OPC UA Server. | | |
| 2434 | OPC UA Subscription Manage the data subscription from OPC UA Server. | | |
| 2436 | UA Simulator Server Create server simulations. | | |
| 2444 | Forwarded Inputs | | |

FIG. 13

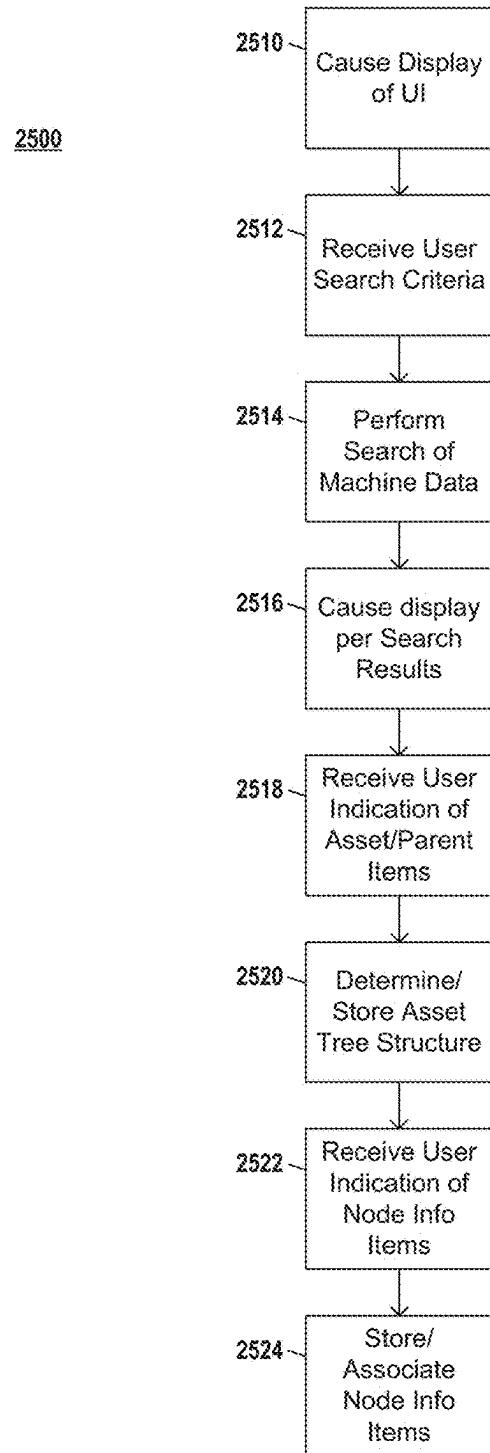
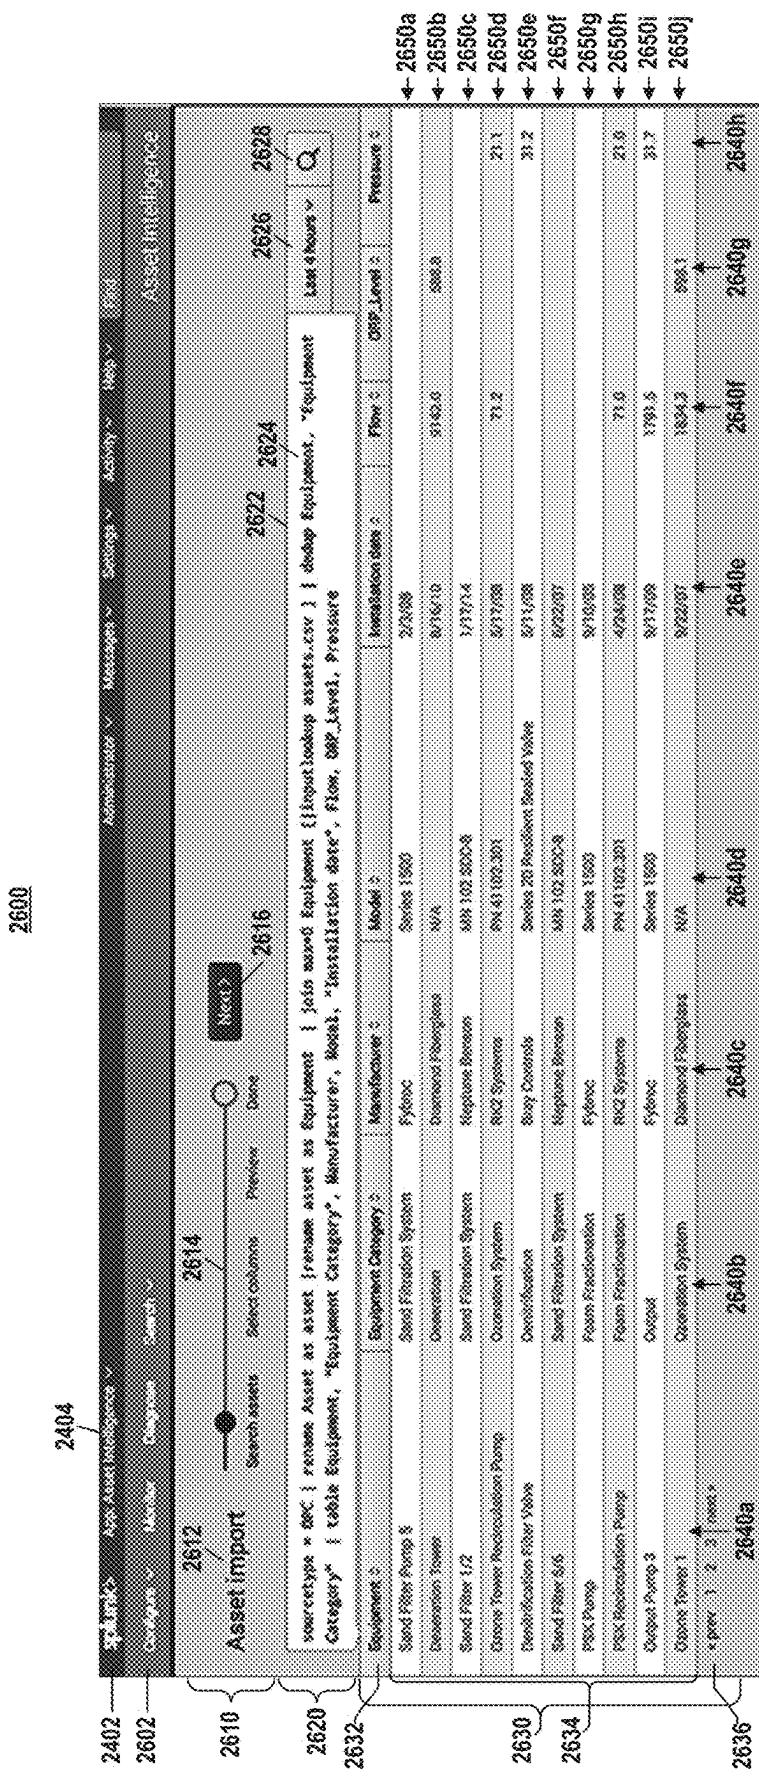
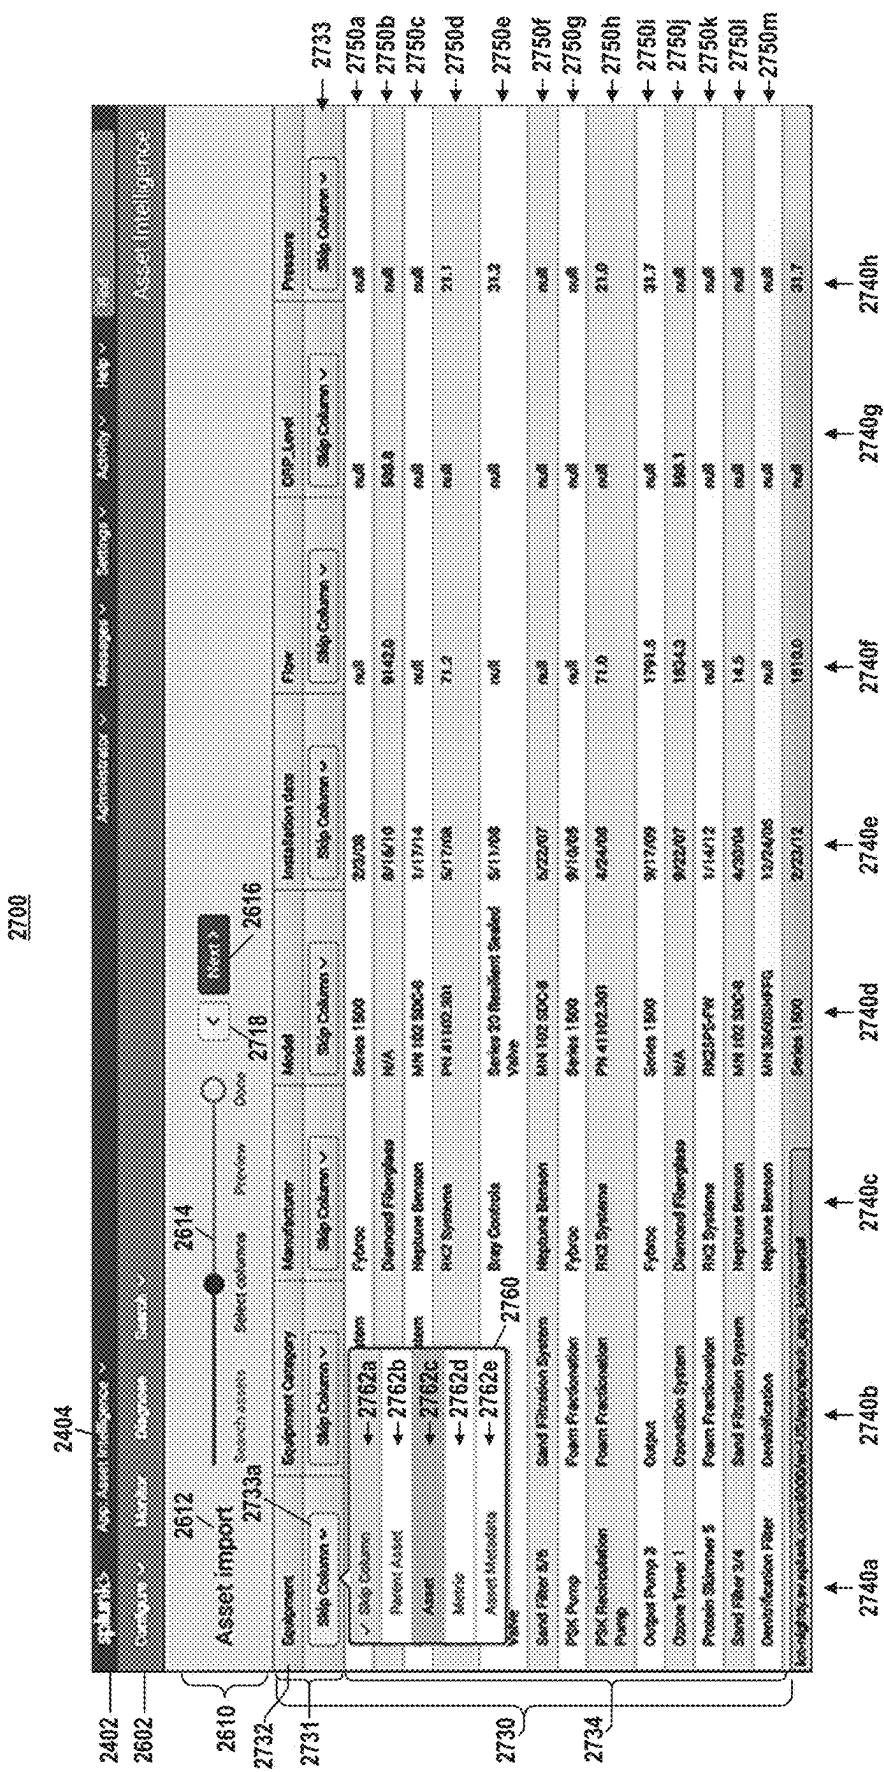


FIG. 14



卷之三



६८

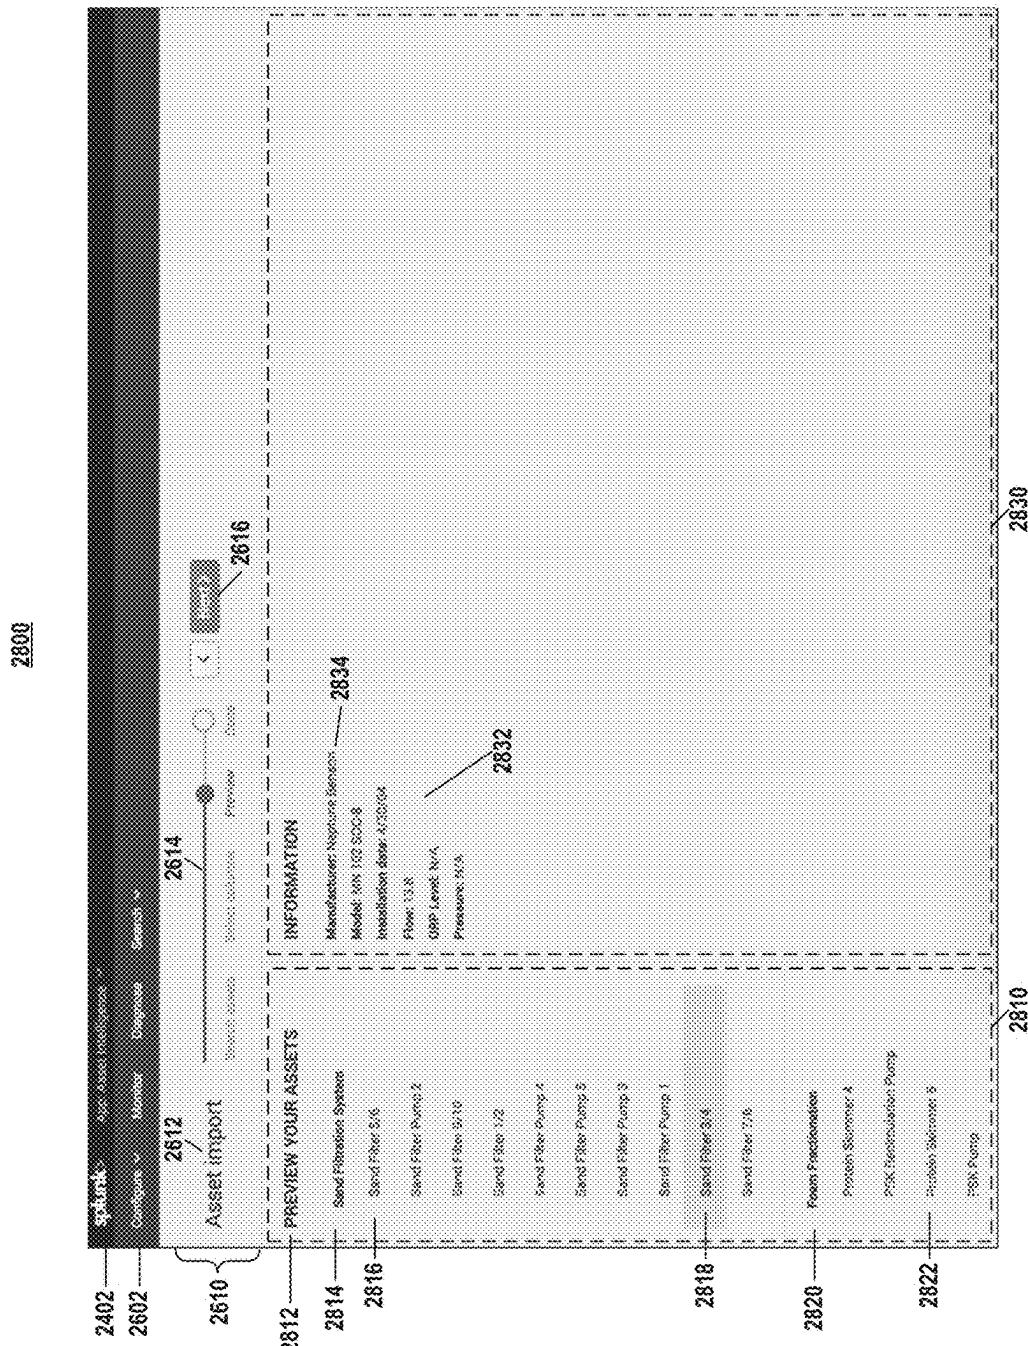


FIG. 17

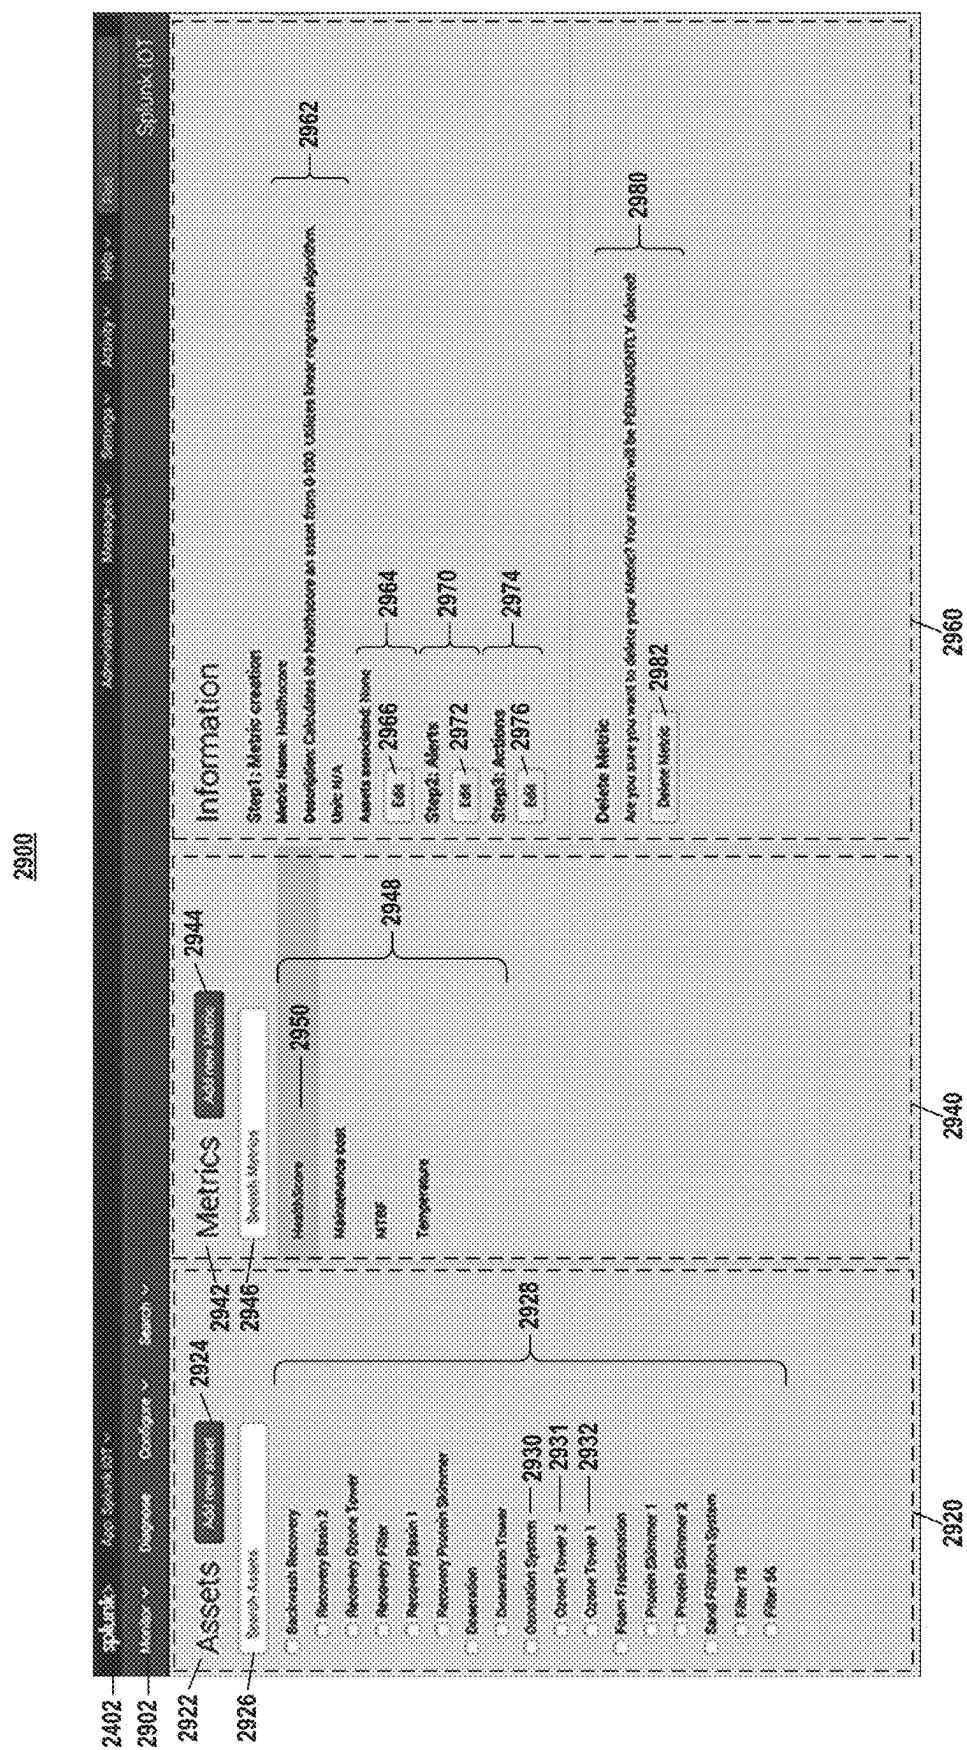


FIG. 18

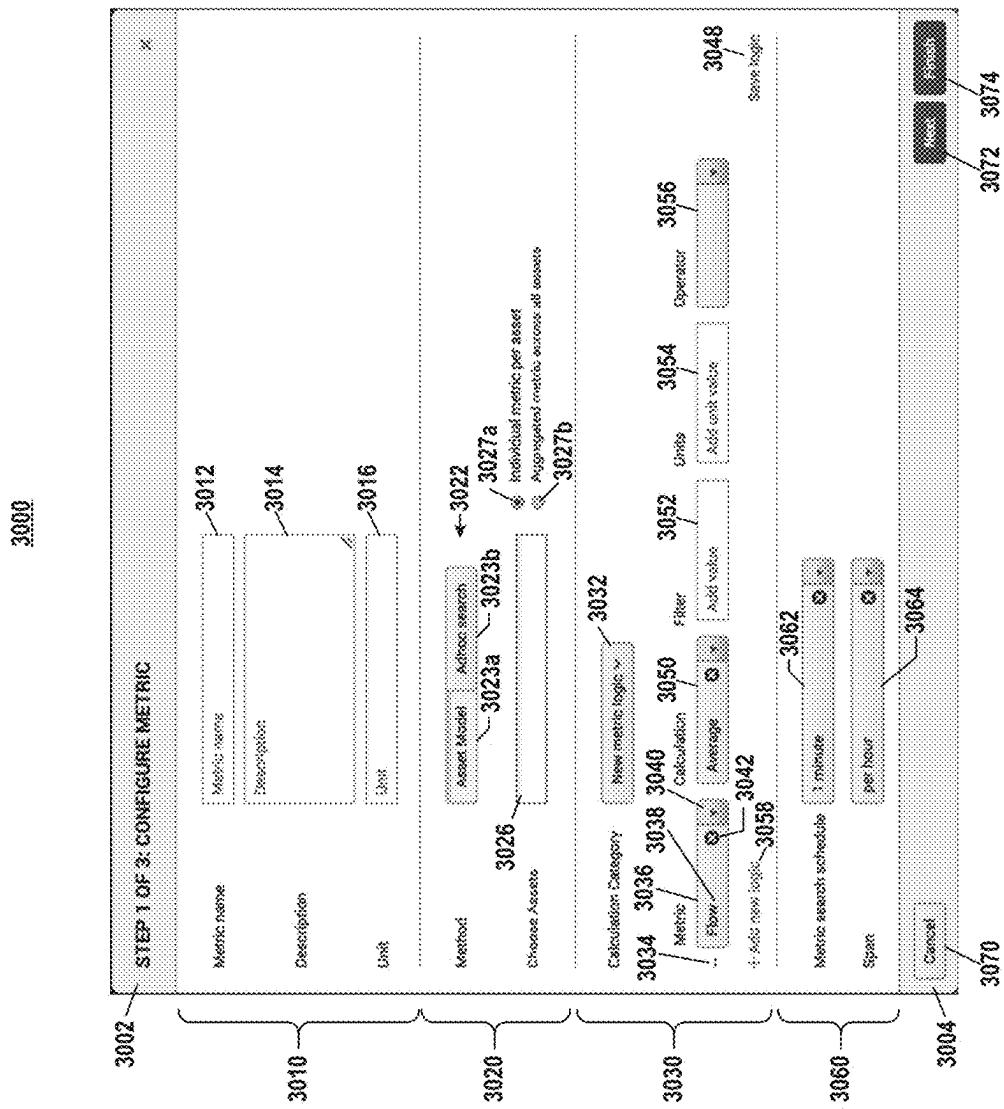


FIG. 19

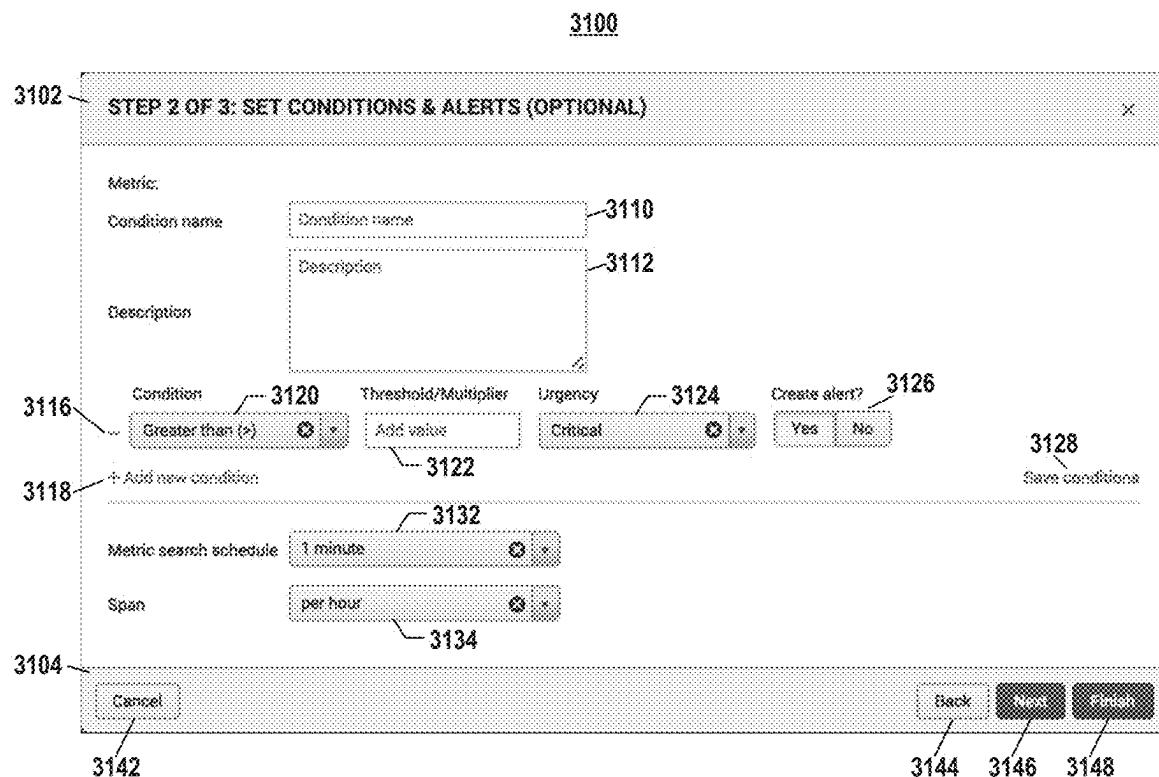


FIG. 20

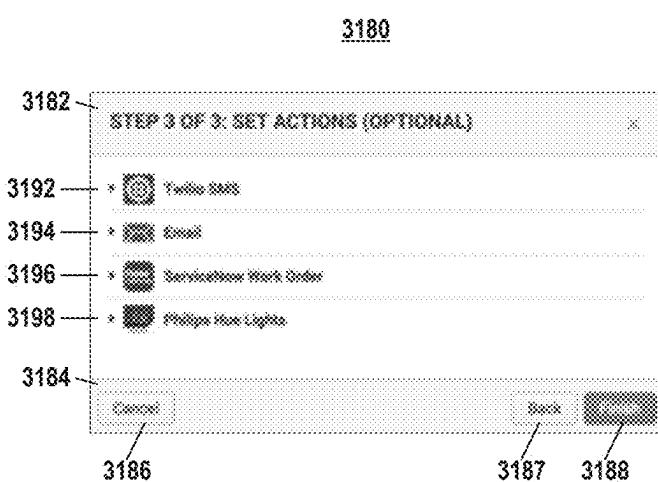


FIG. 21

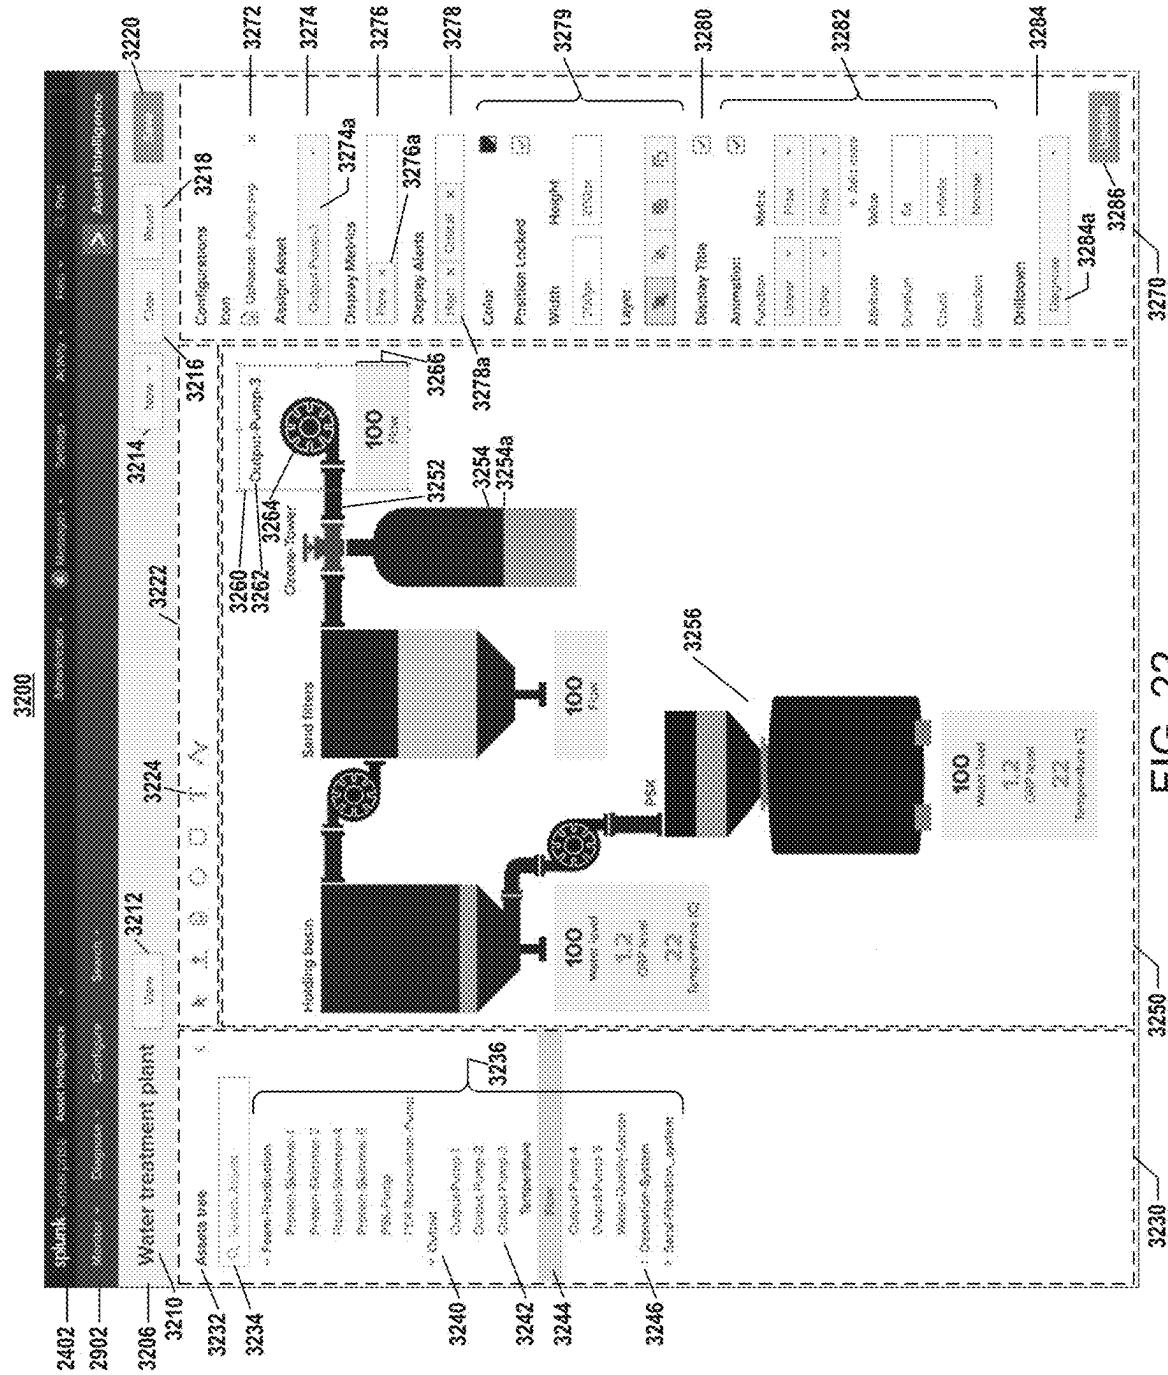


FIG. 22

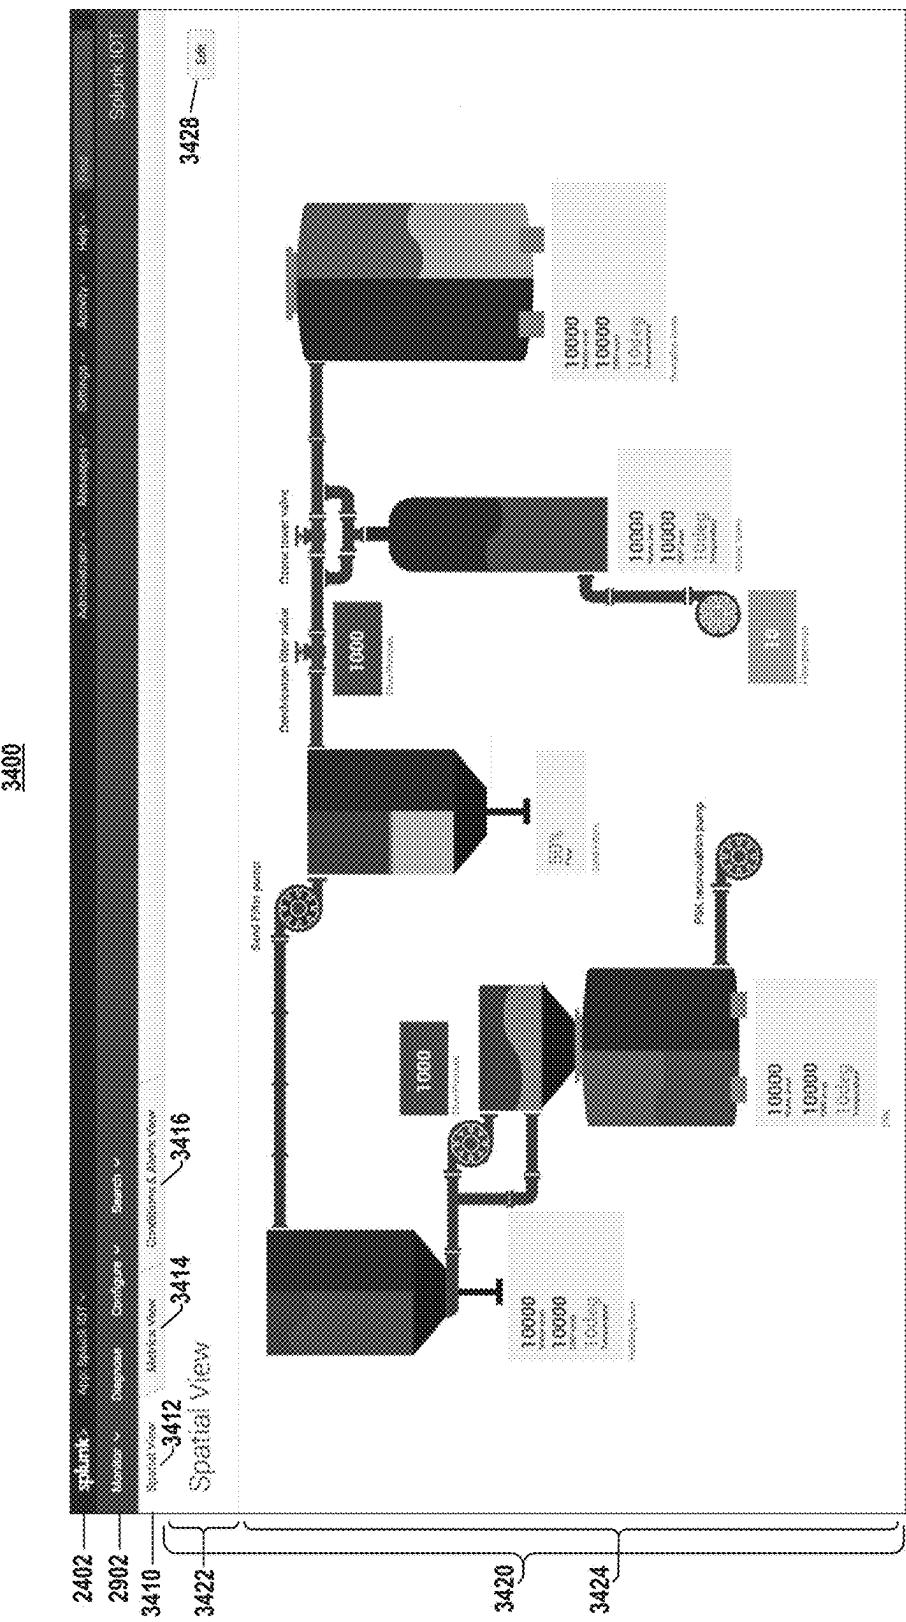


FIG. 23

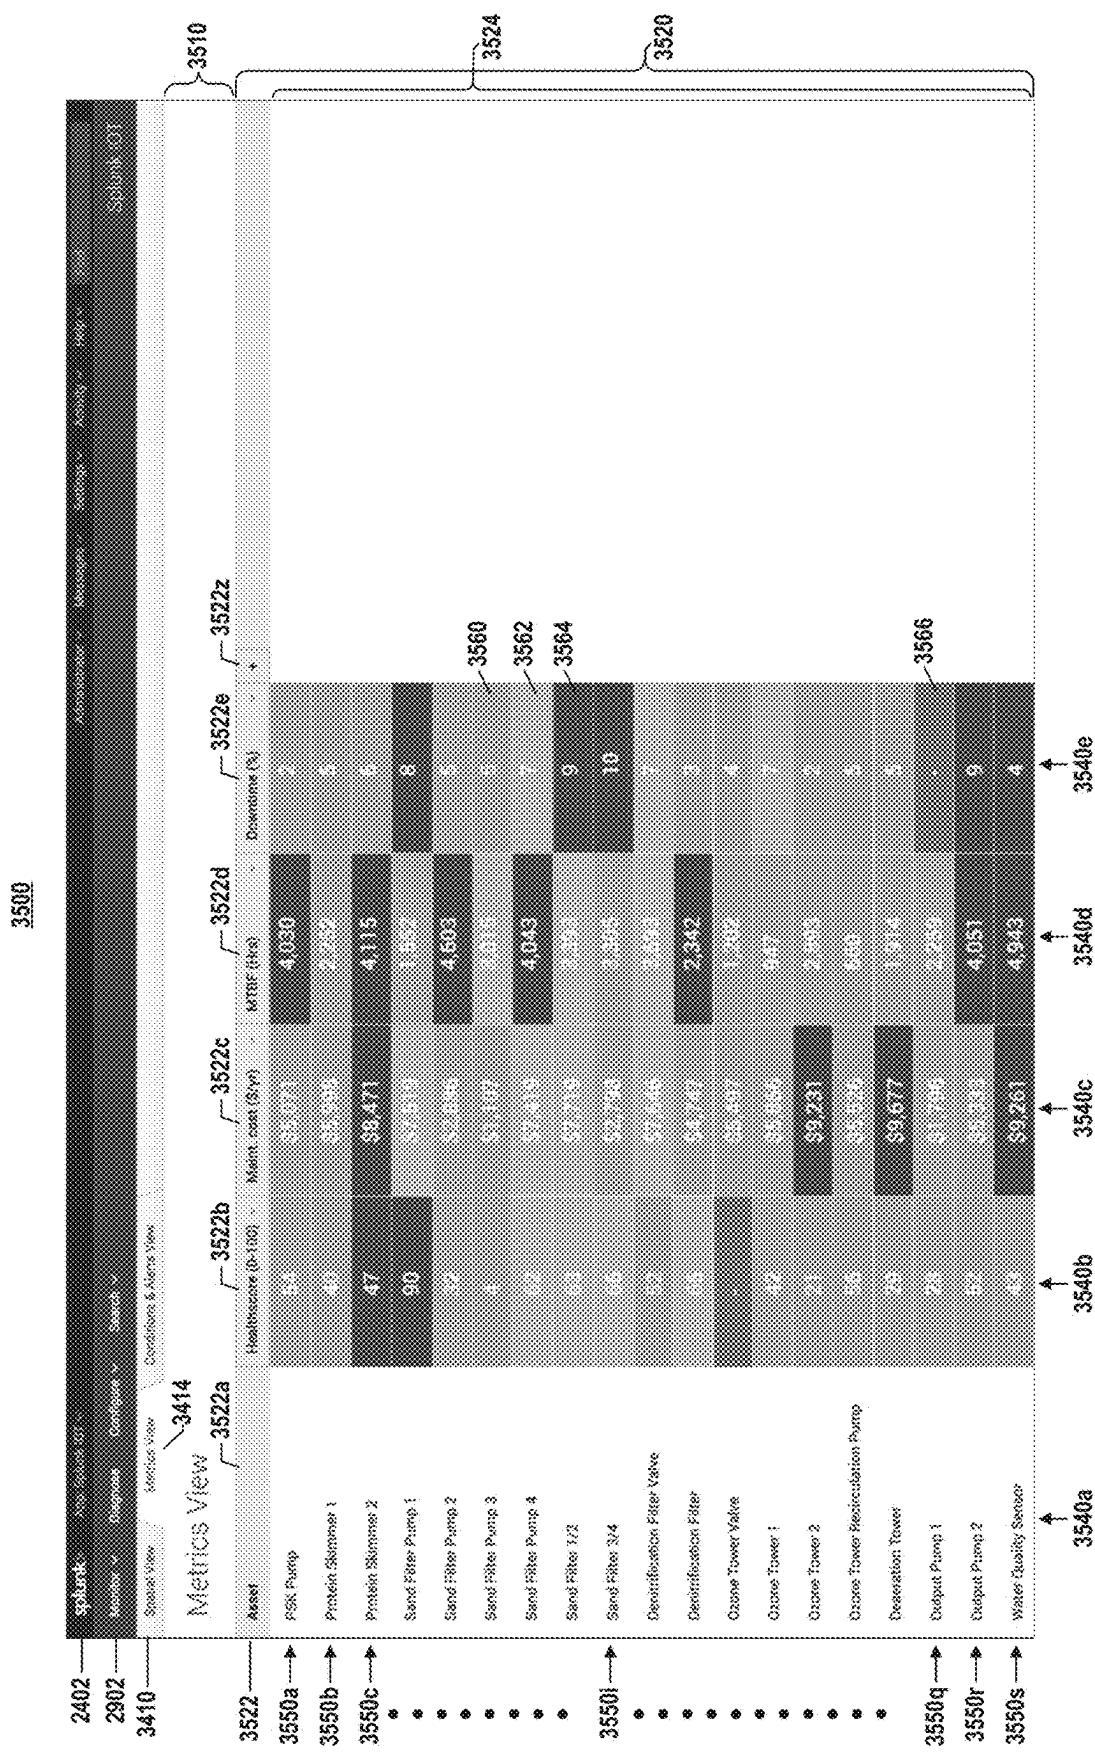


FIG. 24

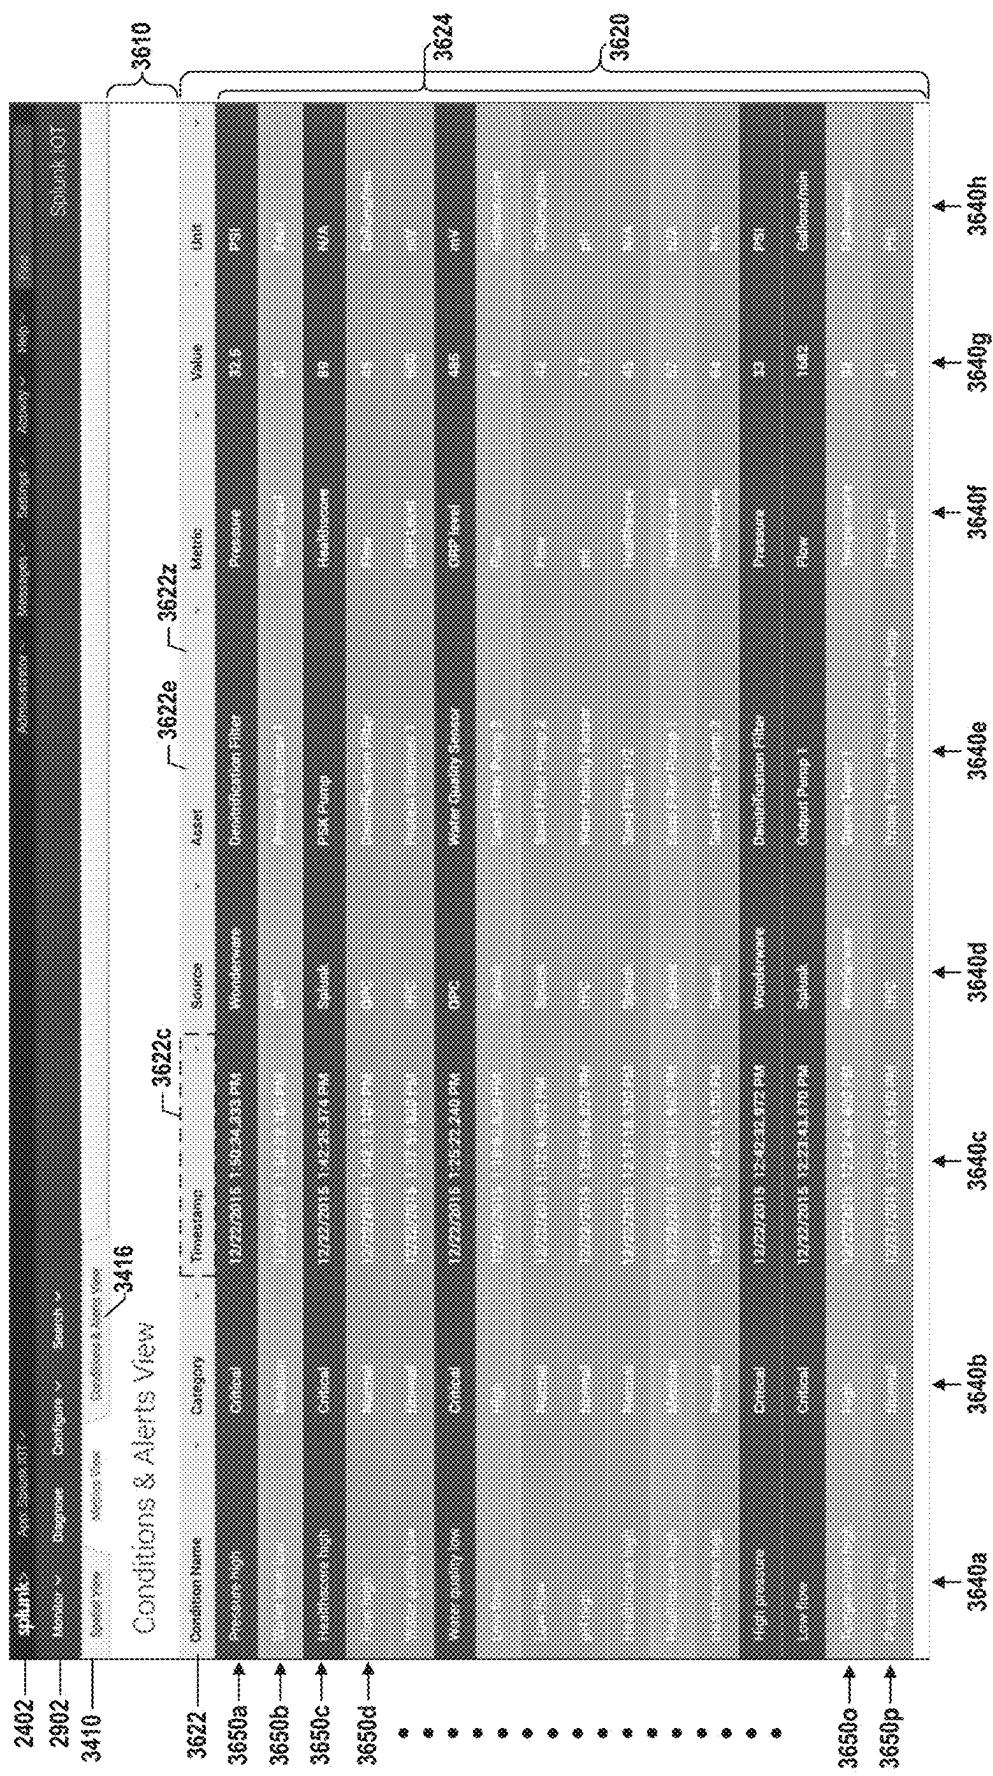


FIG. 25

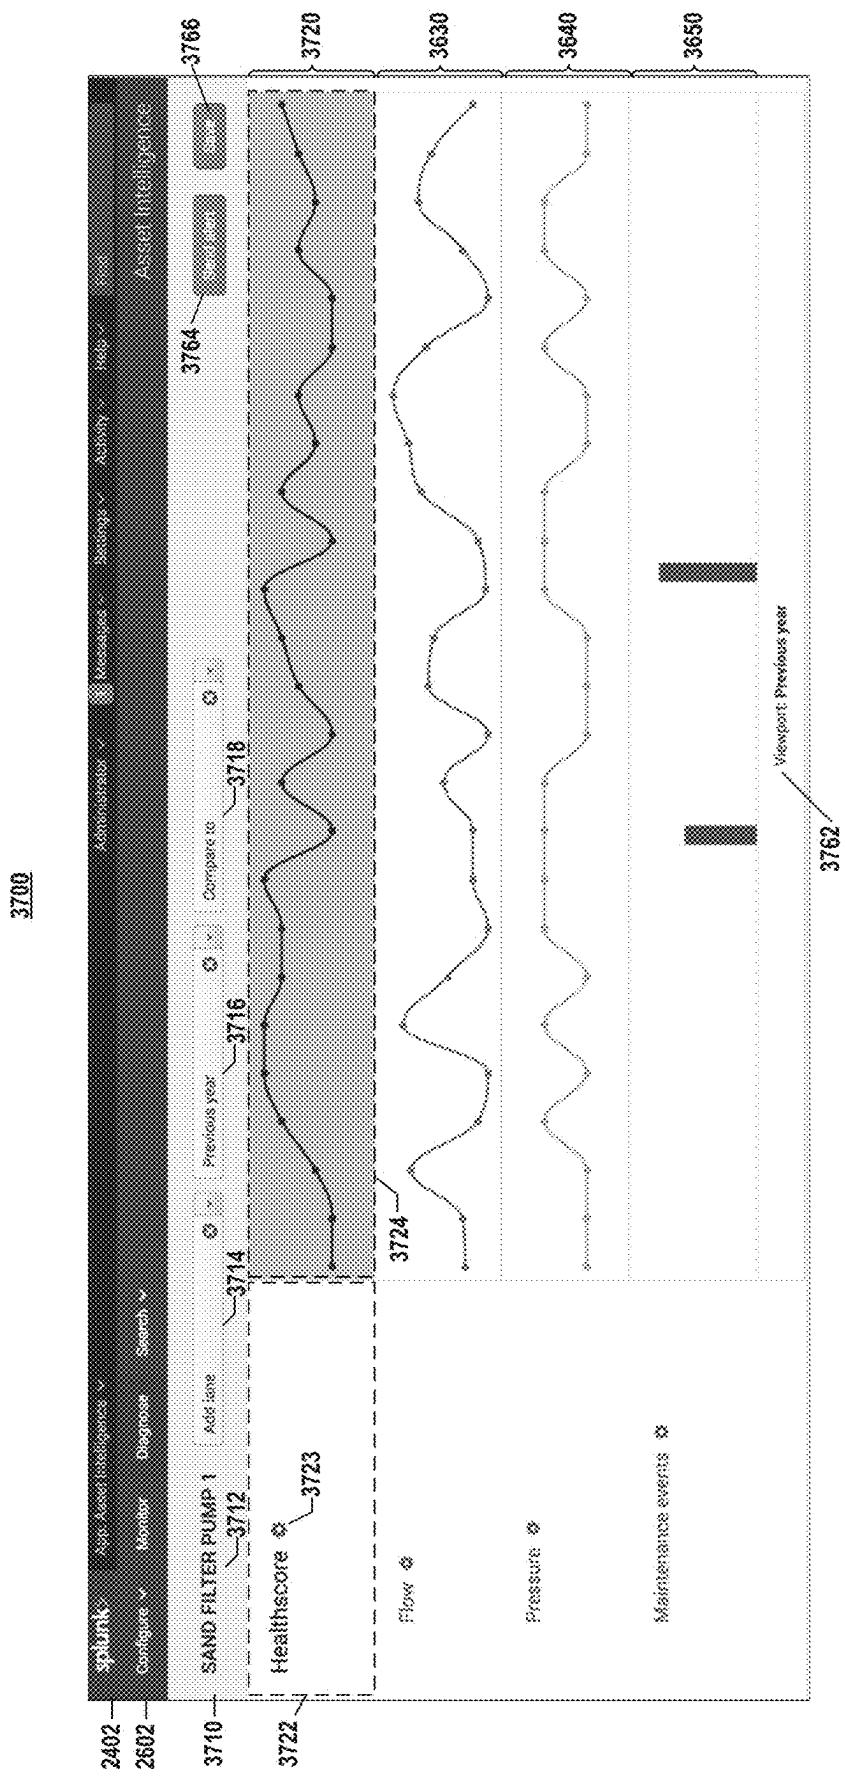


FIG. 26

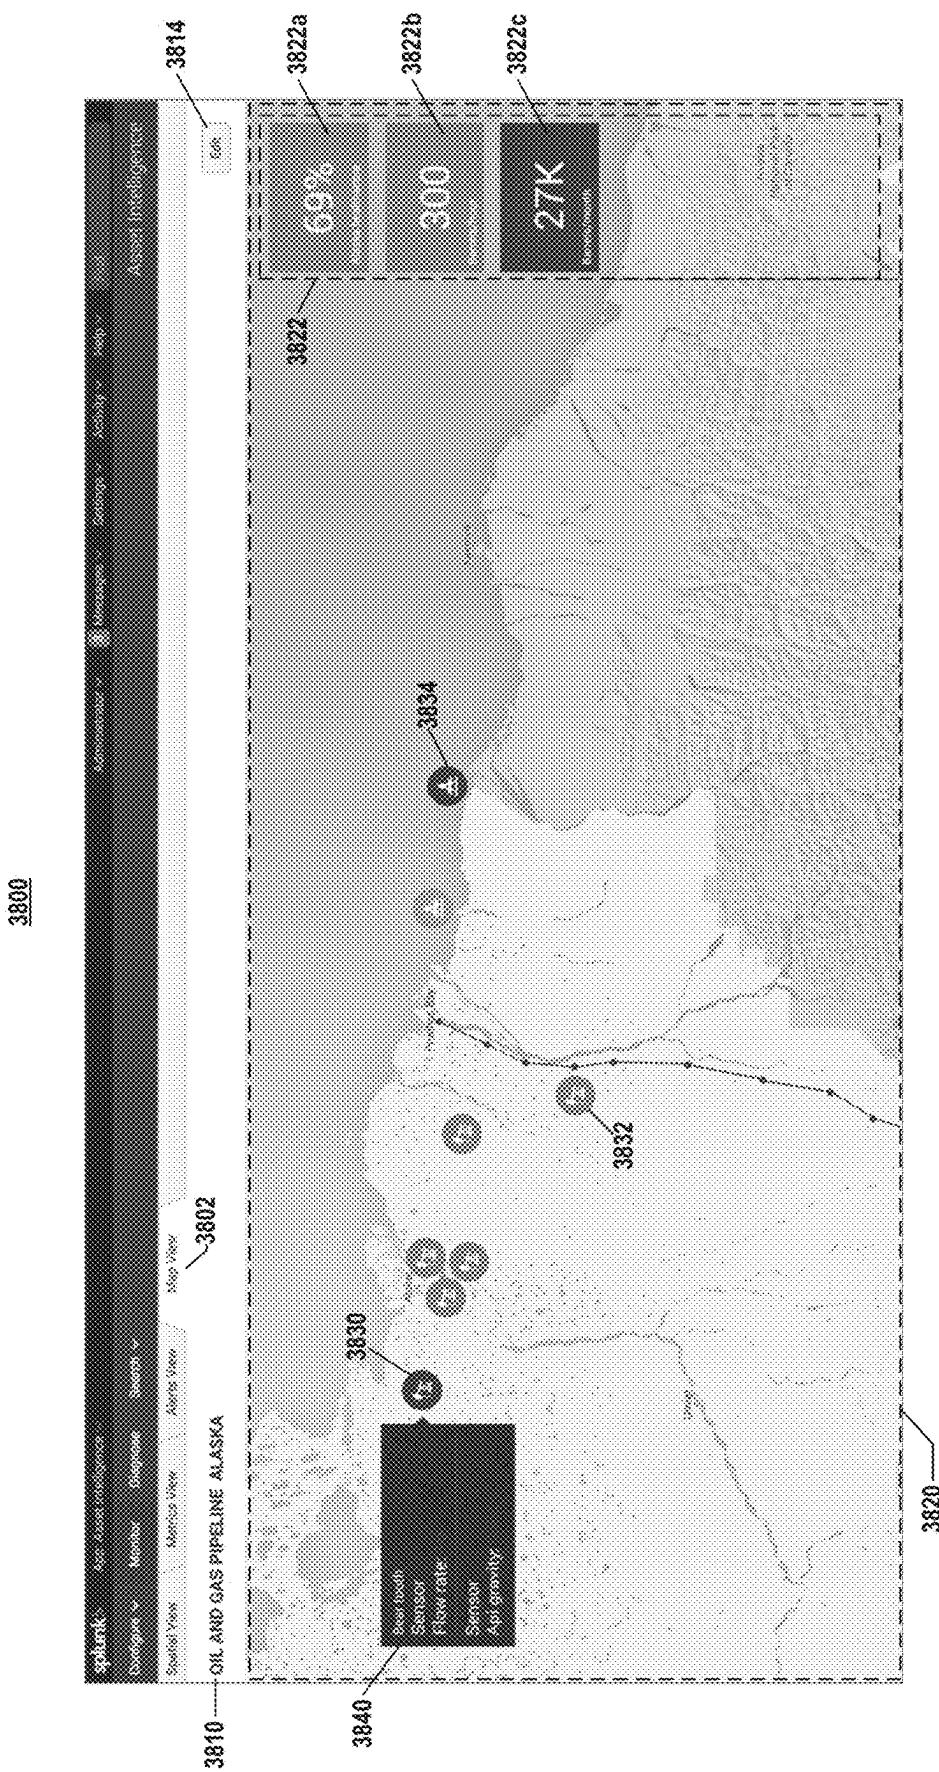


FIG. 27

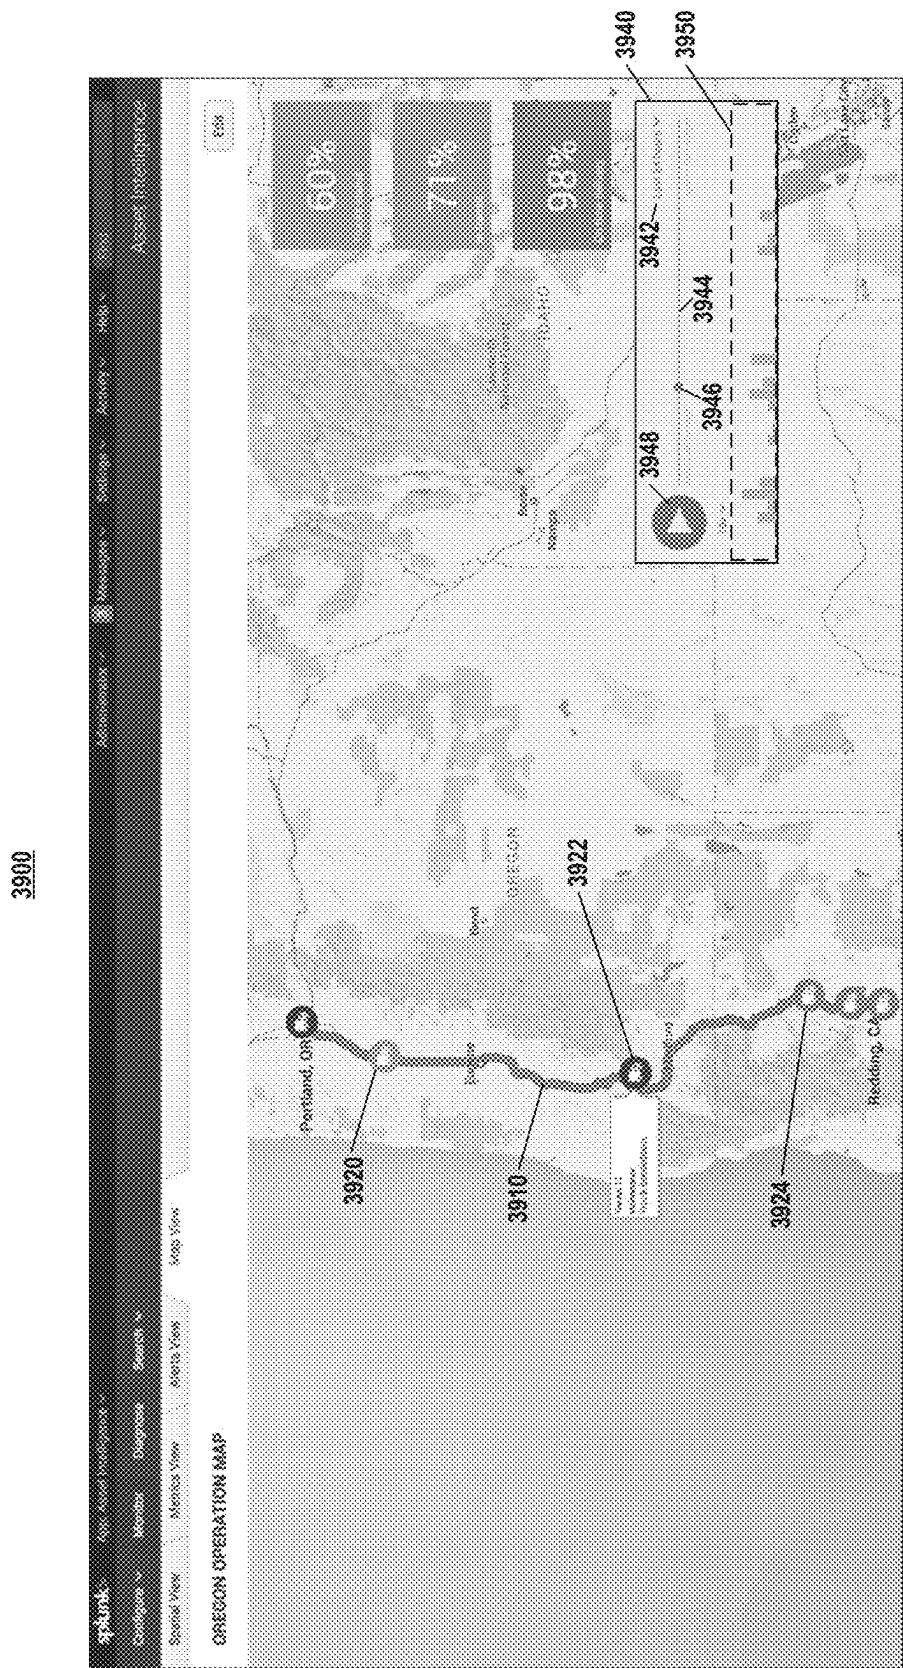


FIG. 28

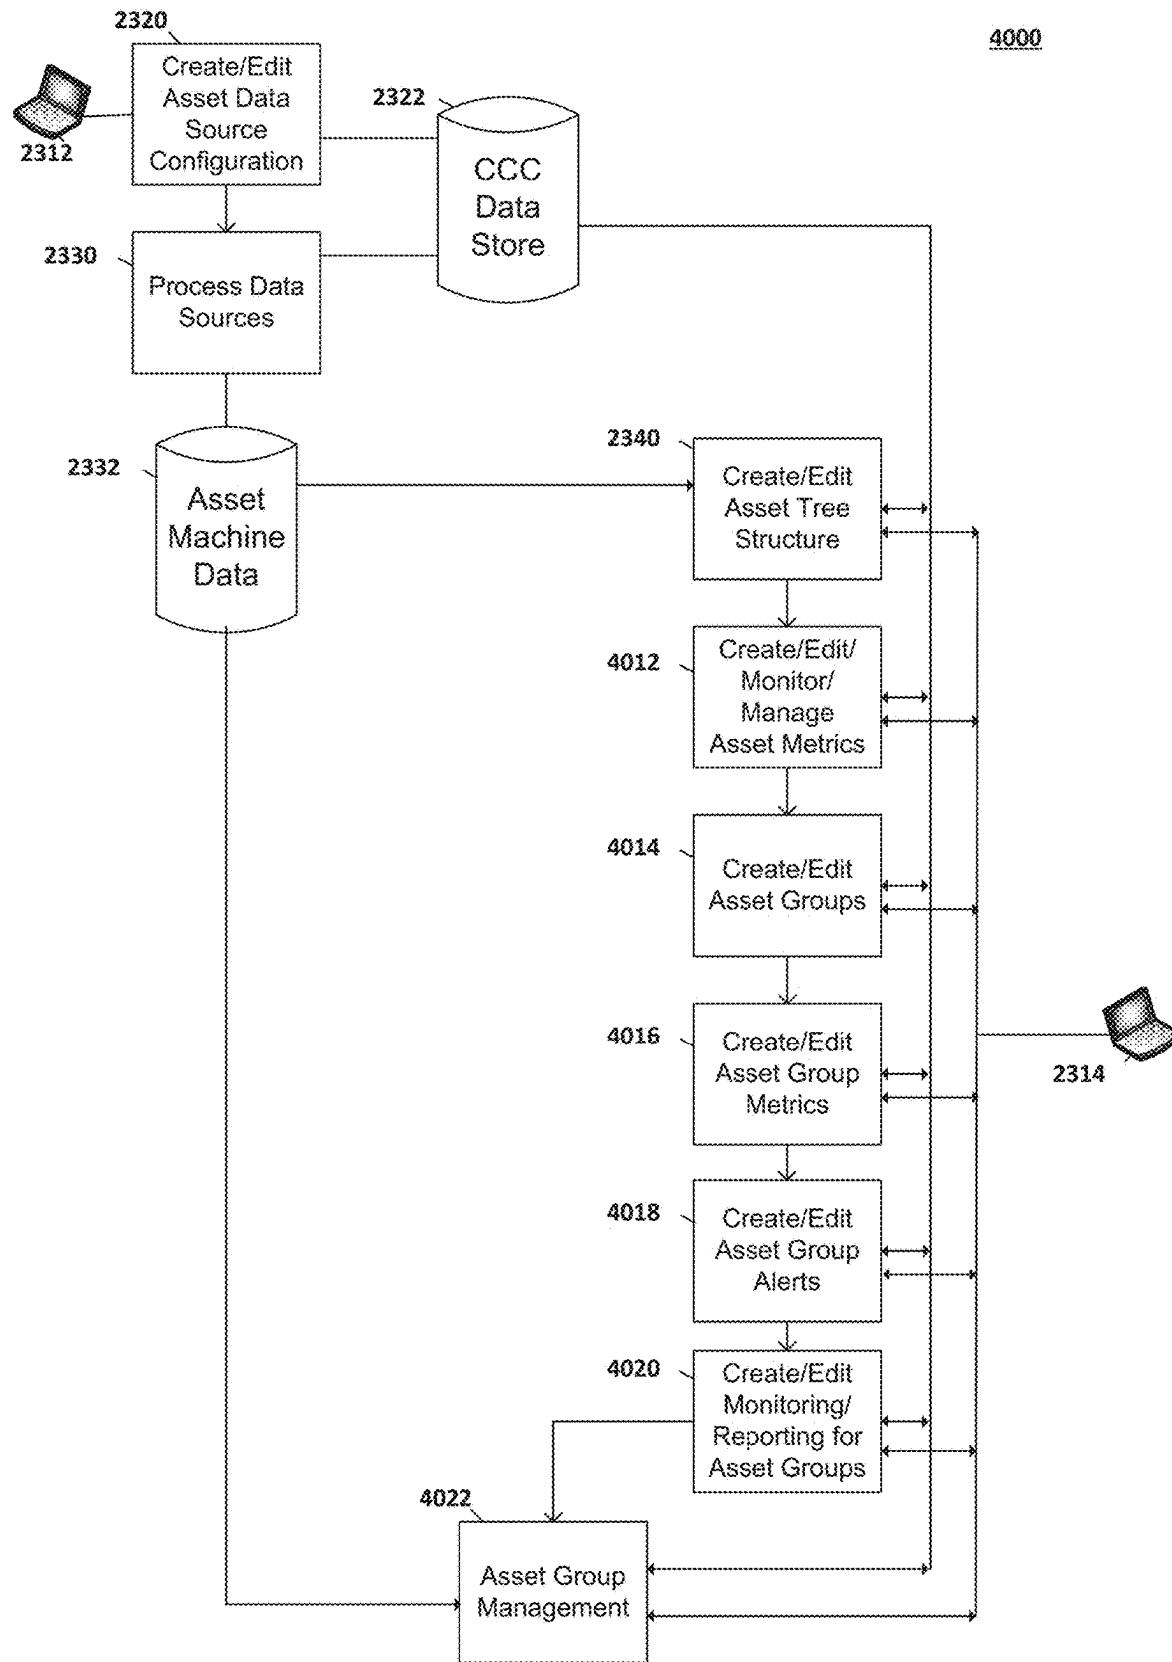


FIG. 29

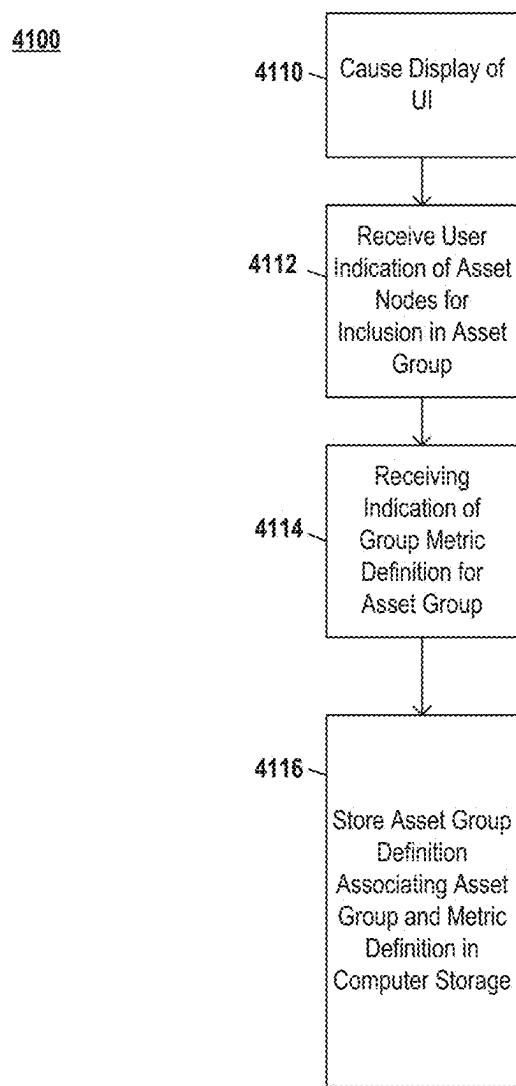


FIG. 30A

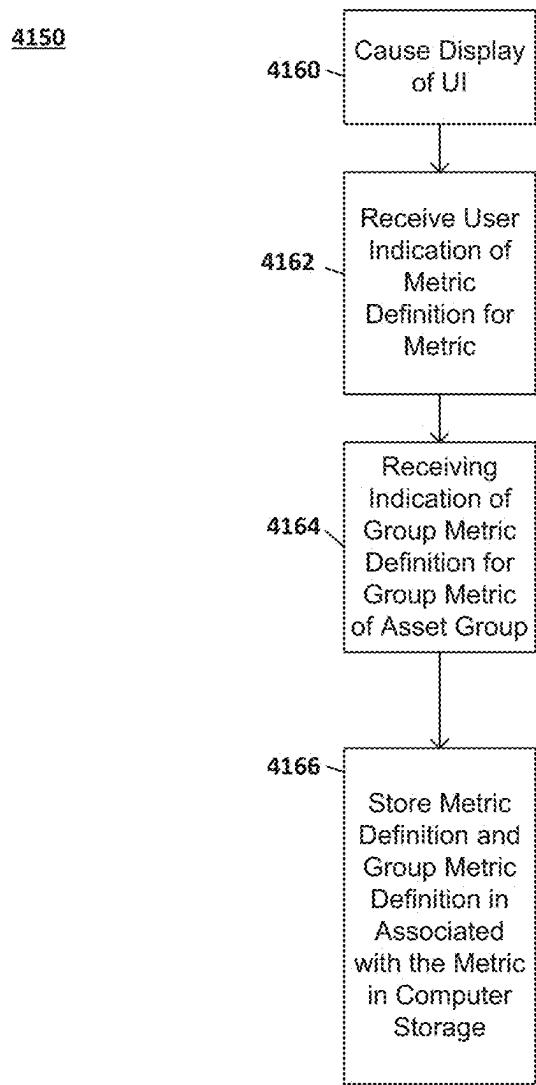


FIG. 30B

4200

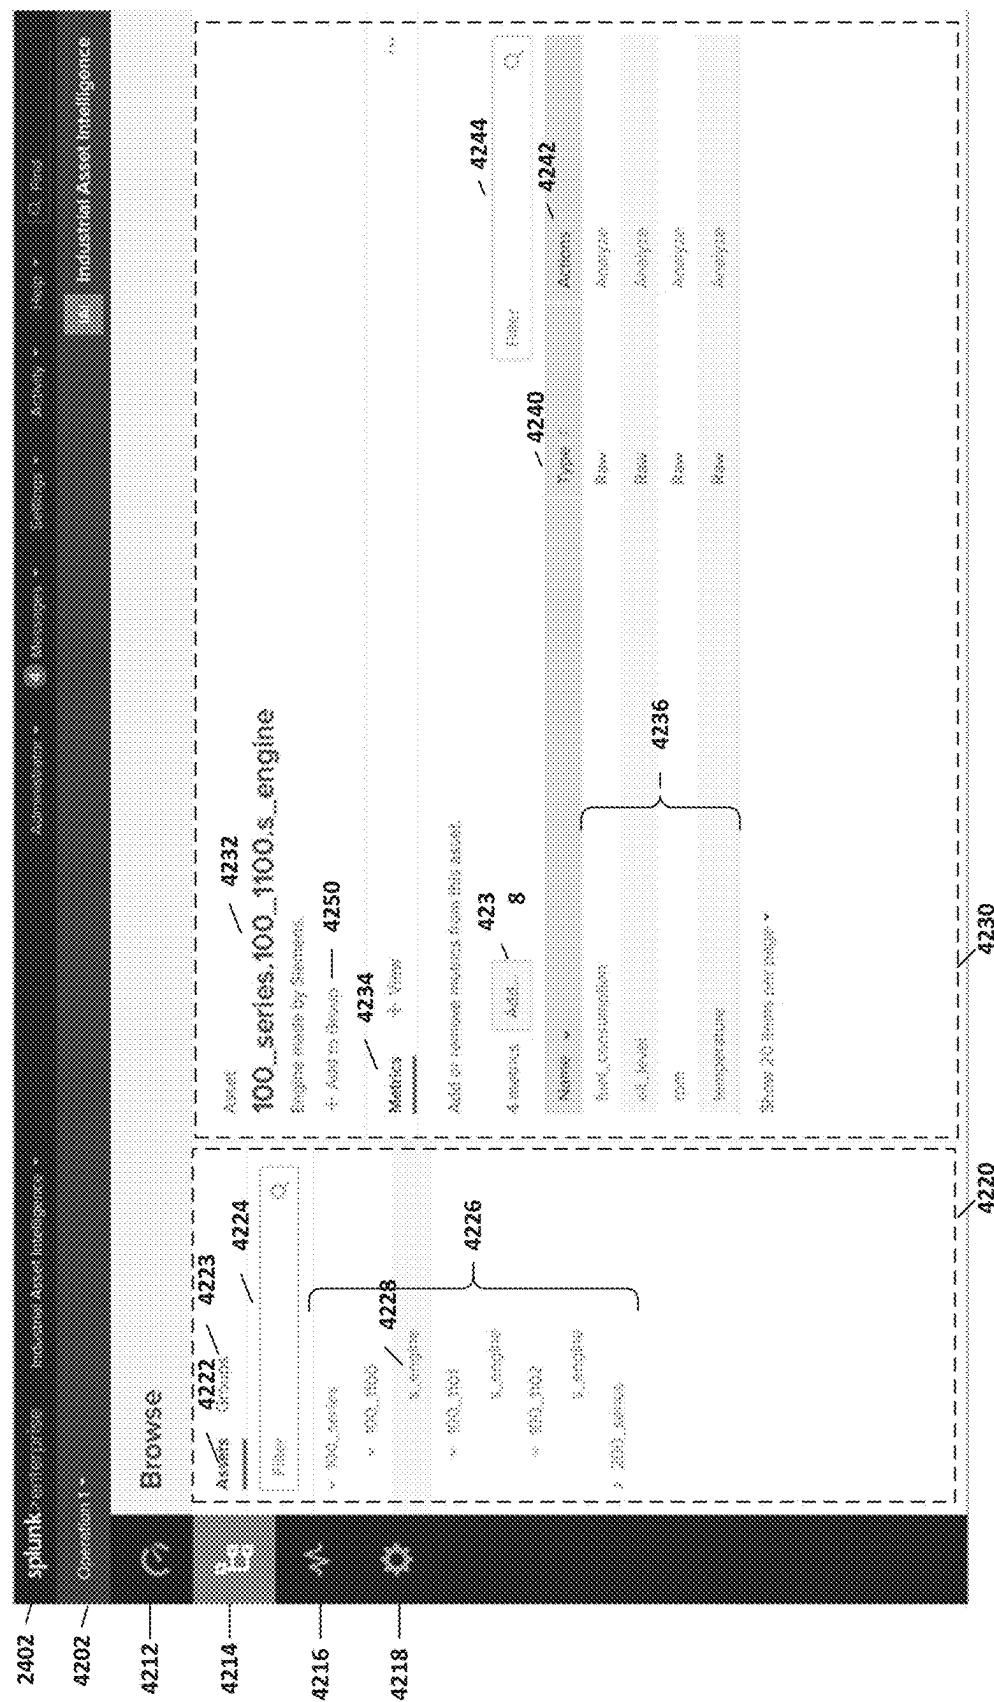


FIG. 31

4309

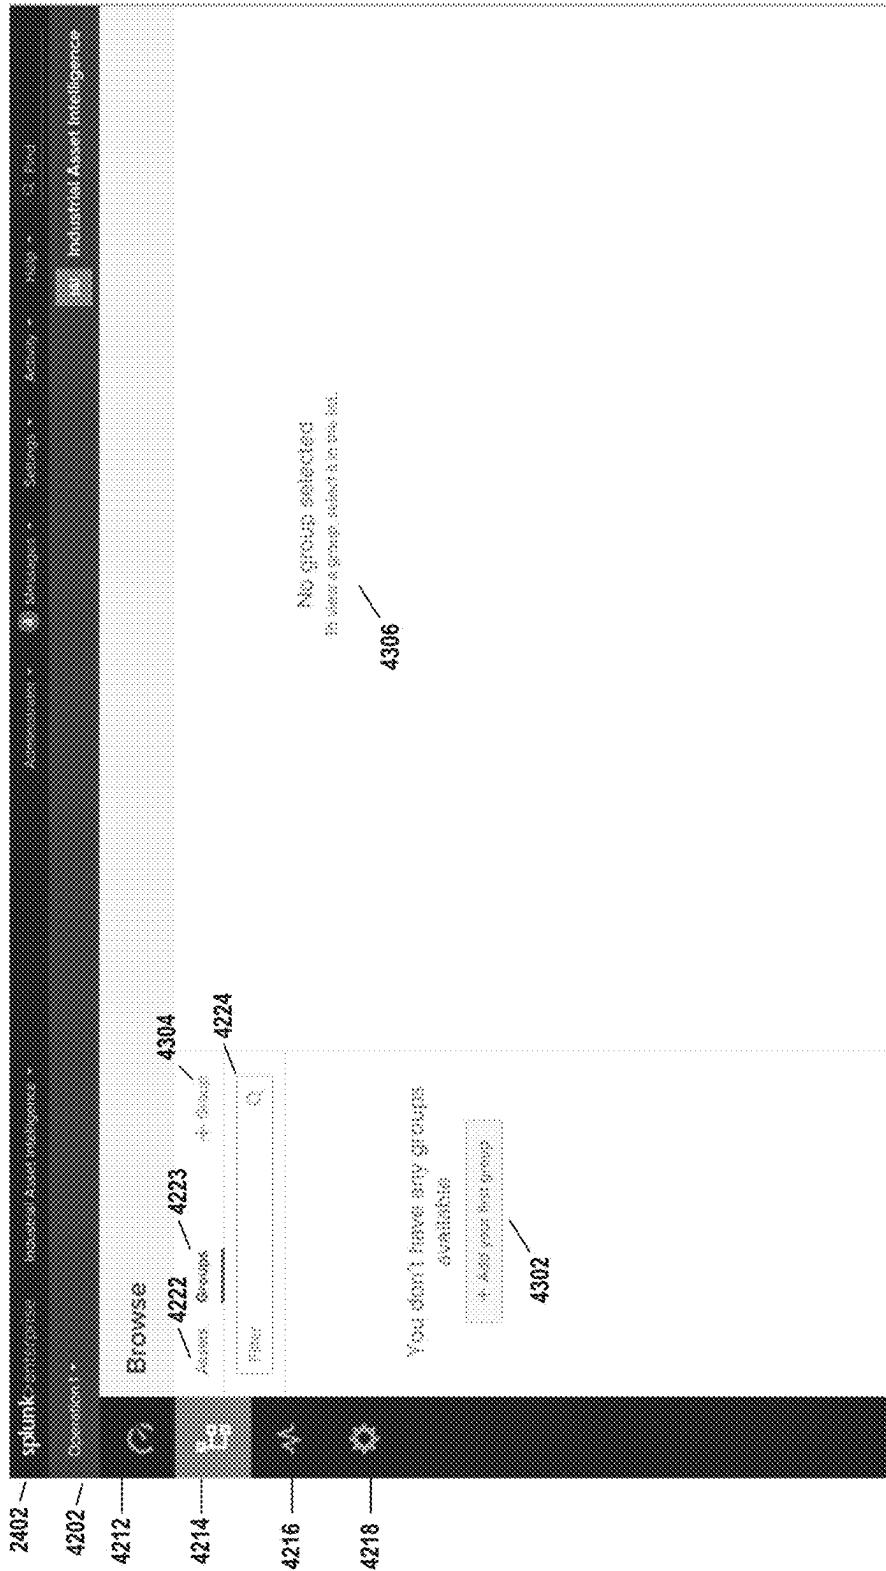


FIG. 32

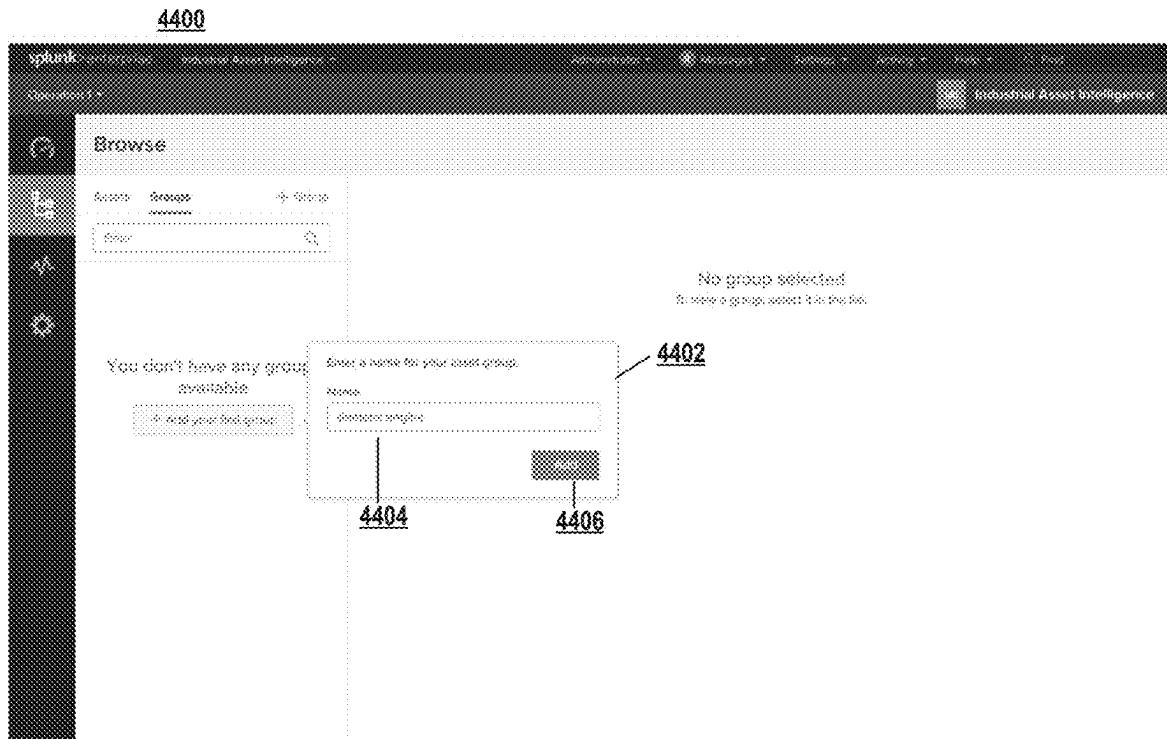


FIG. 33

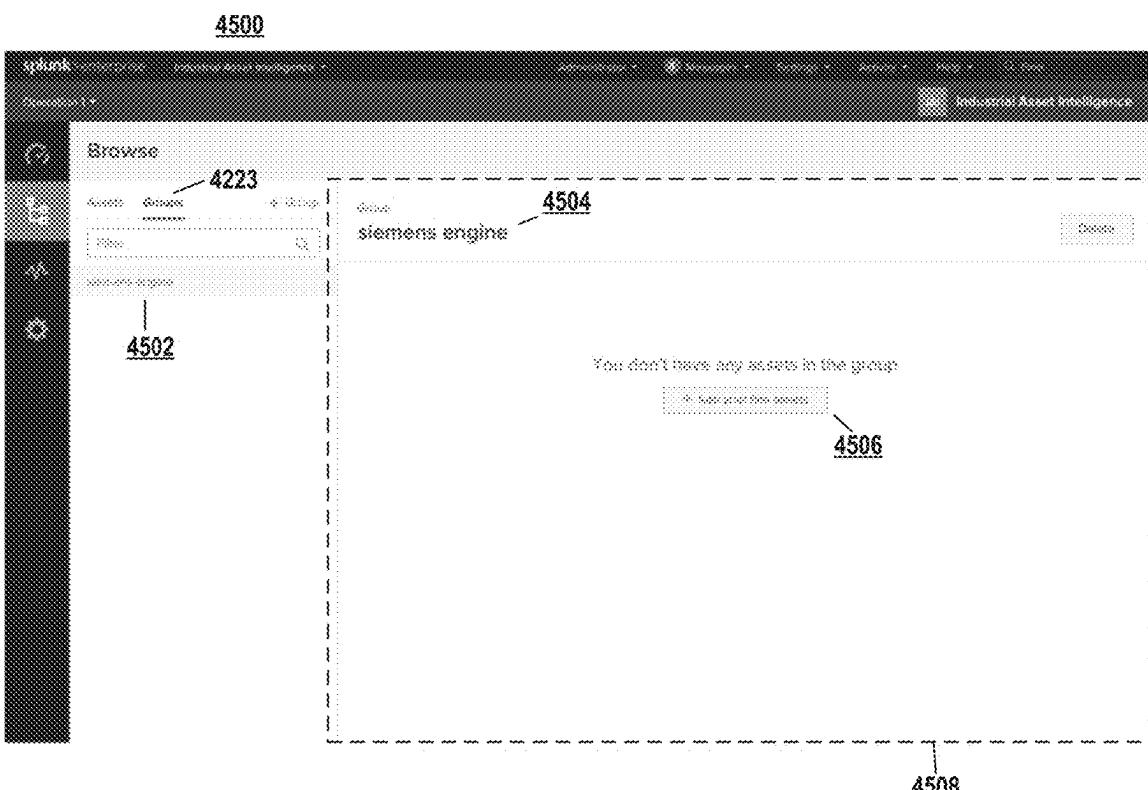


FIG. 34

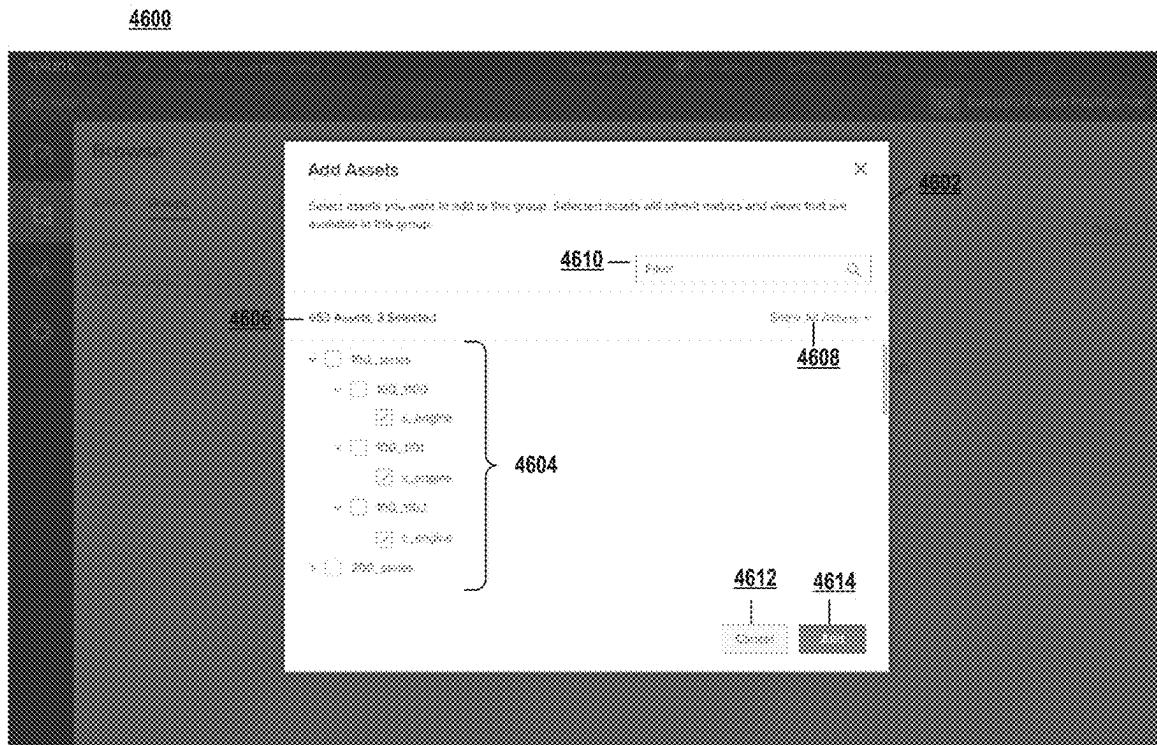


FIG. 35

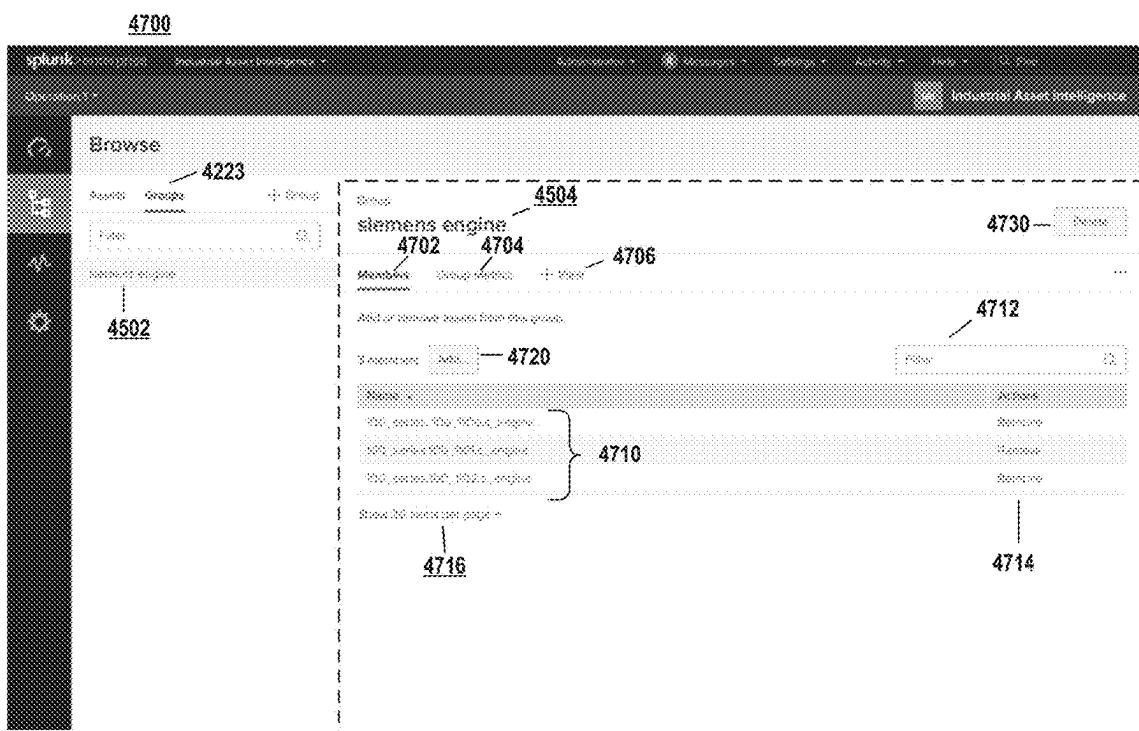


FIG. 36

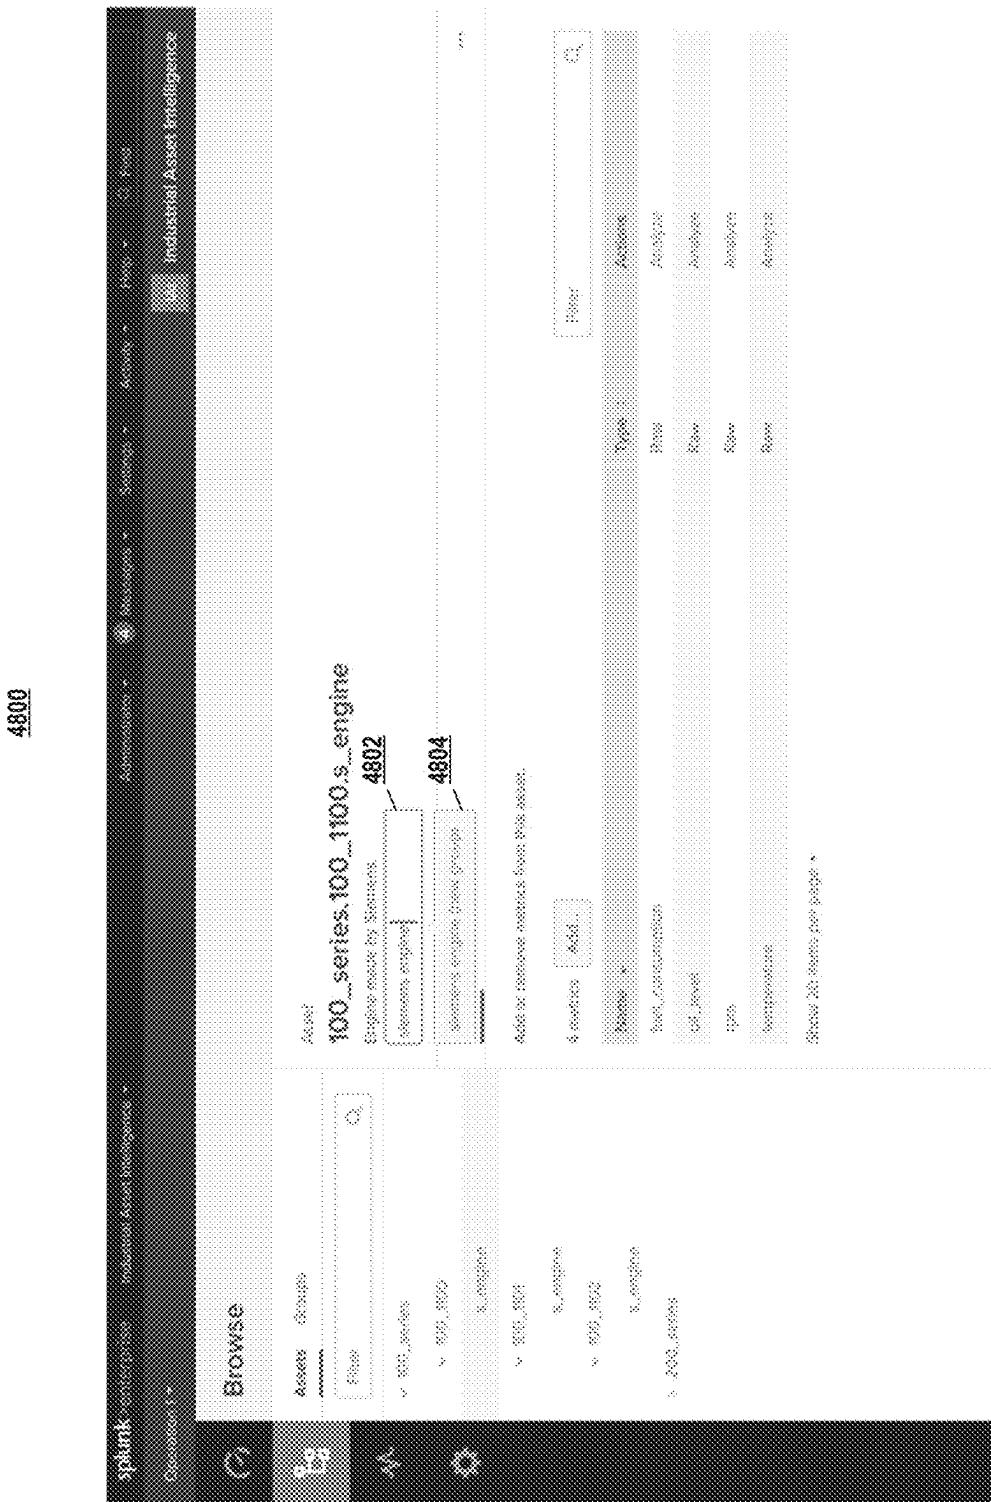


FIG. 37

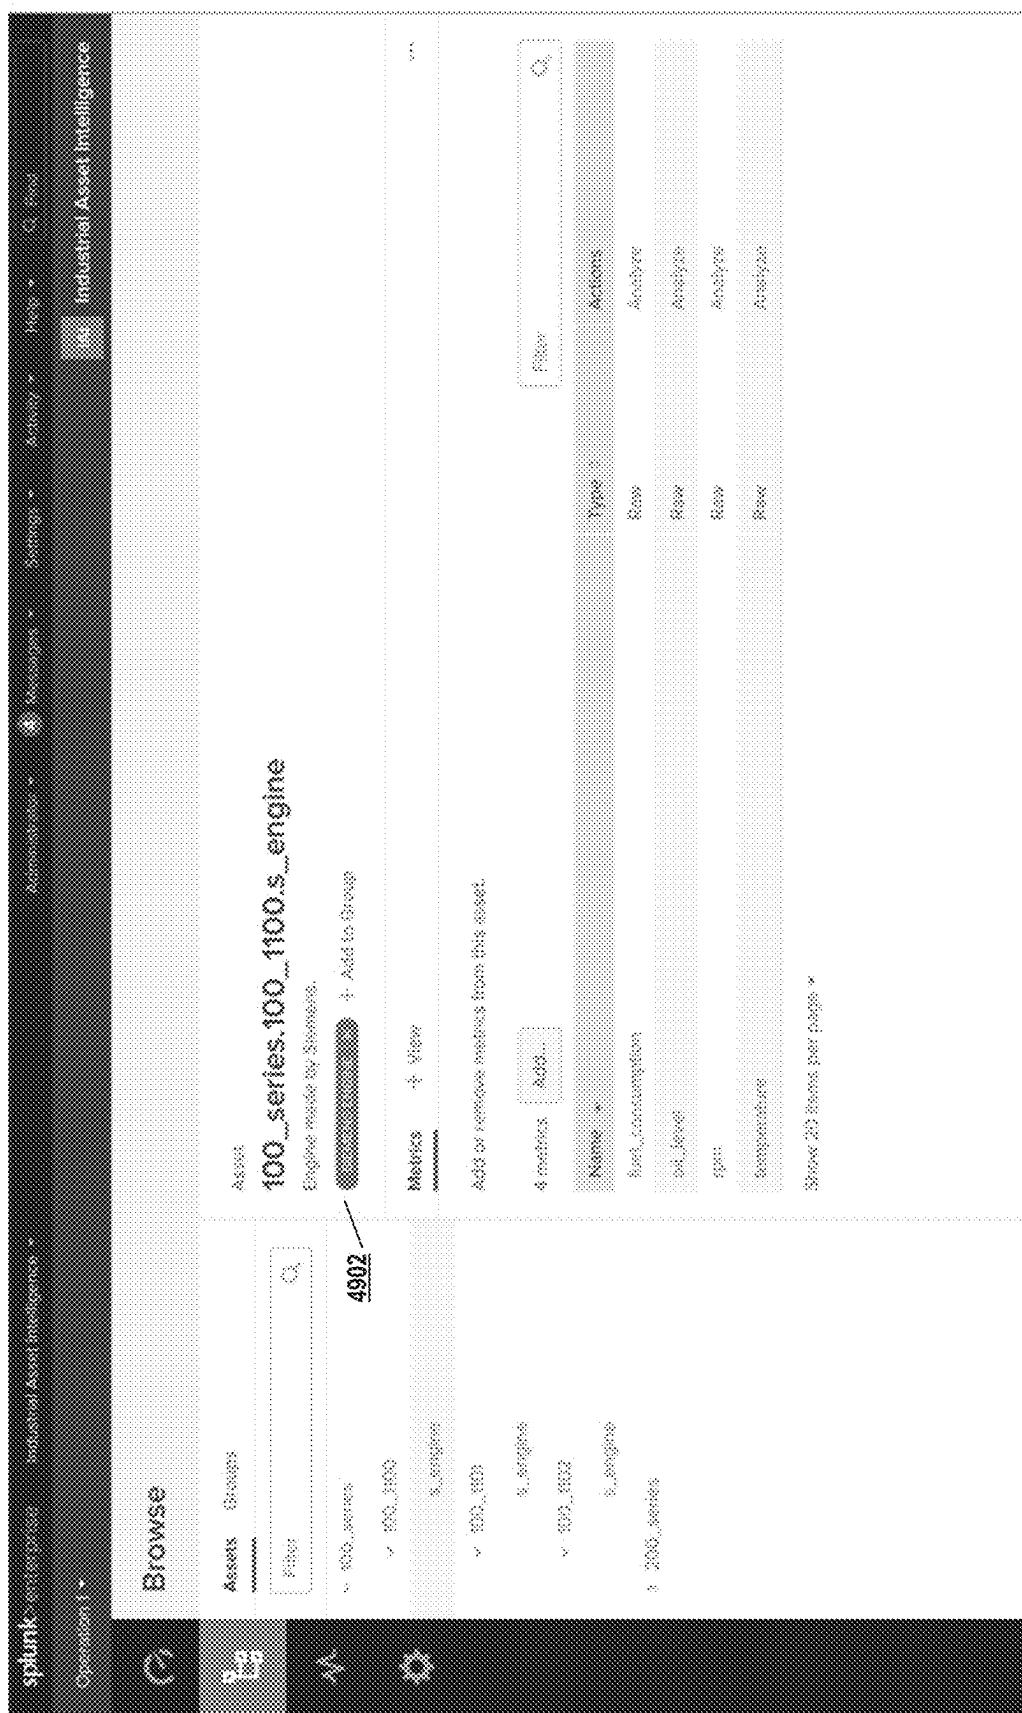


FIG. 38

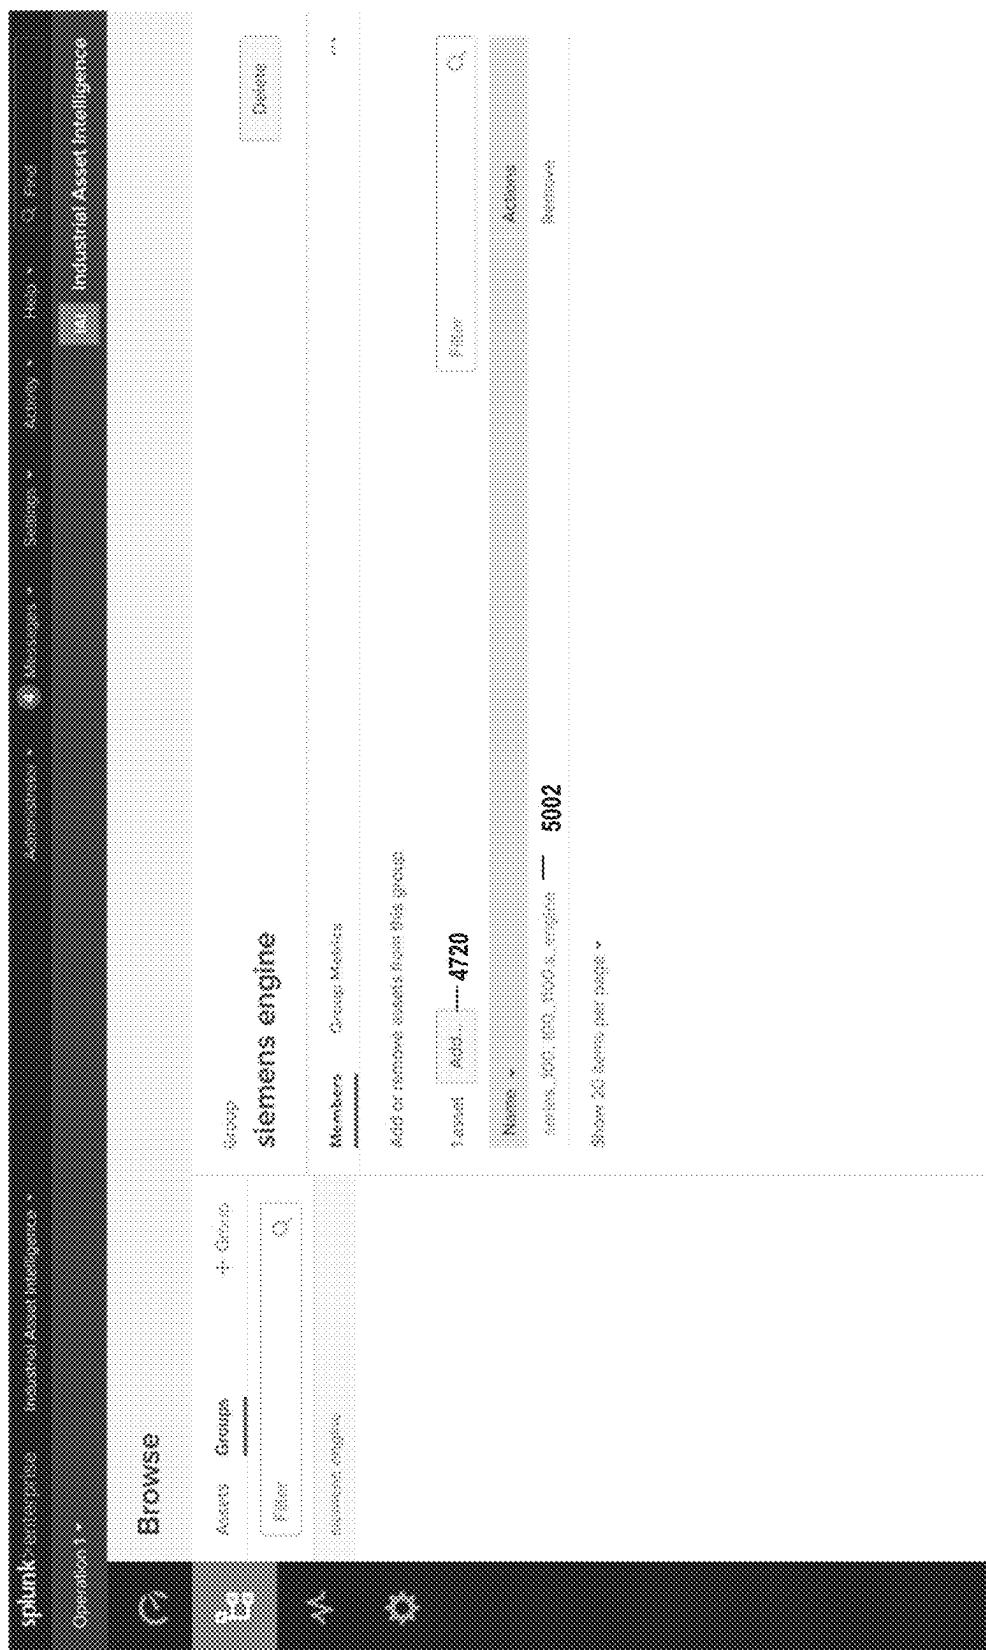


FIG. 39

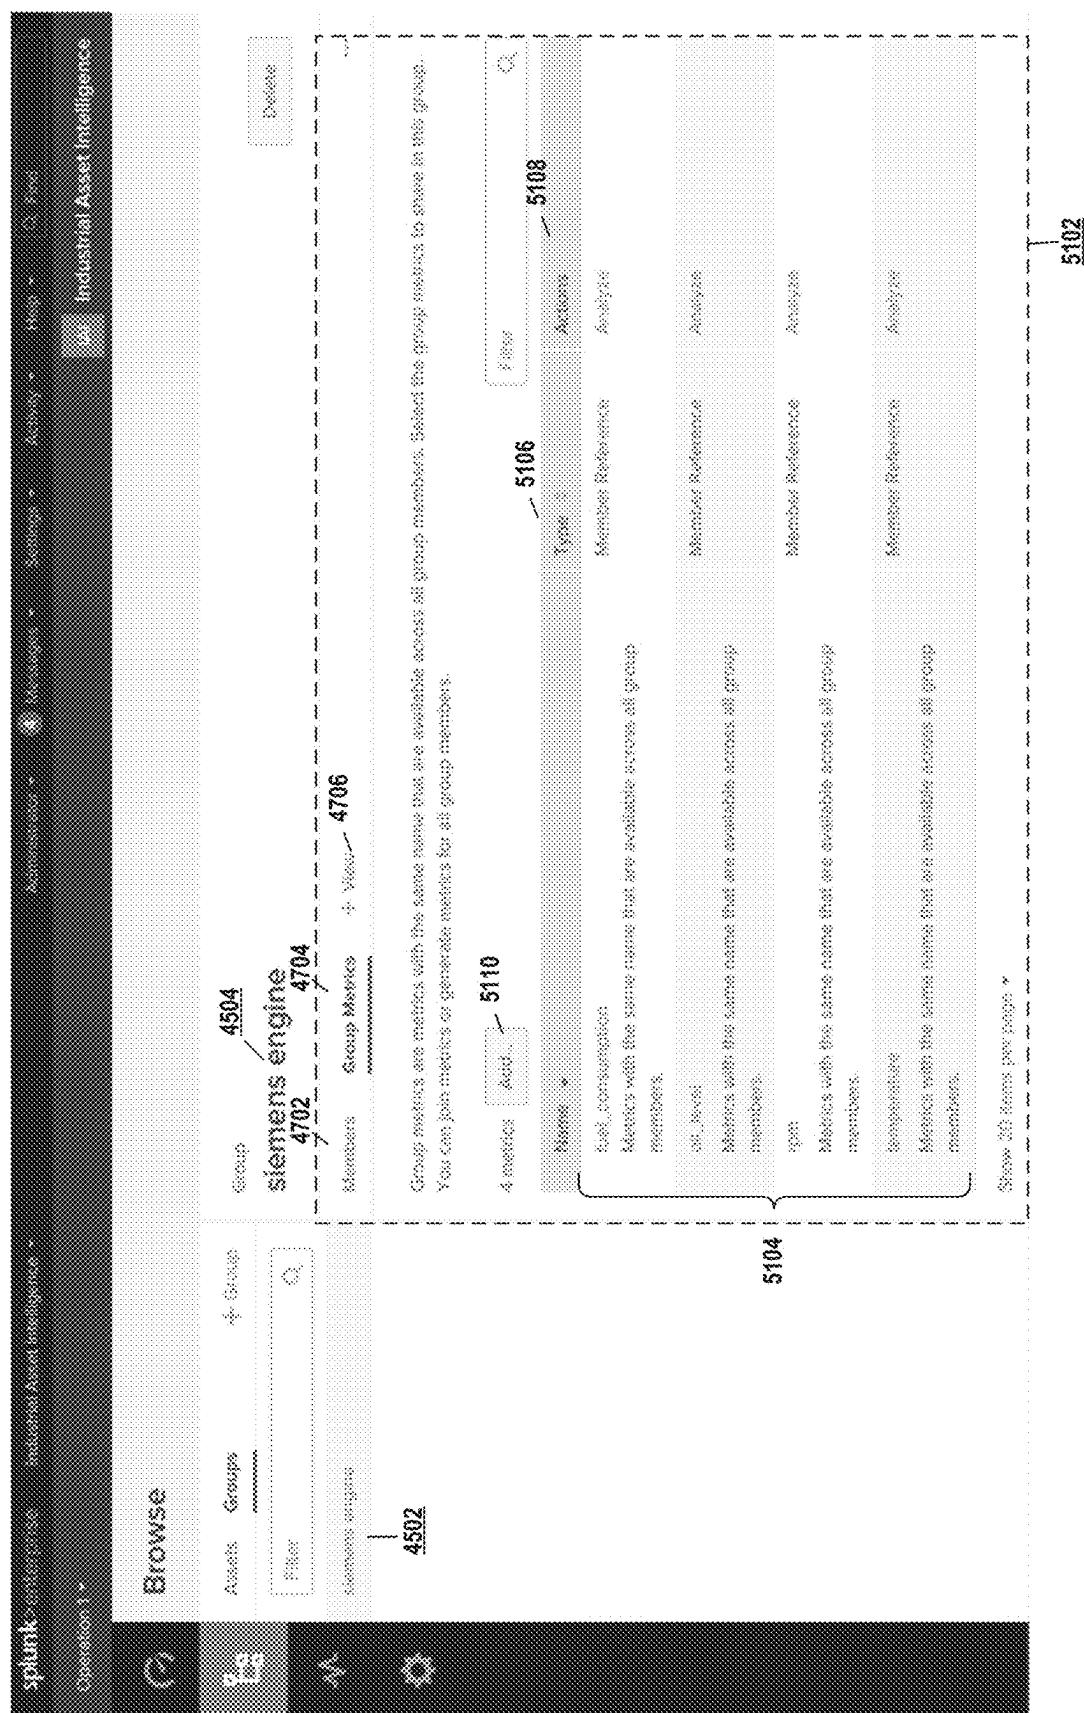


FIG. 40

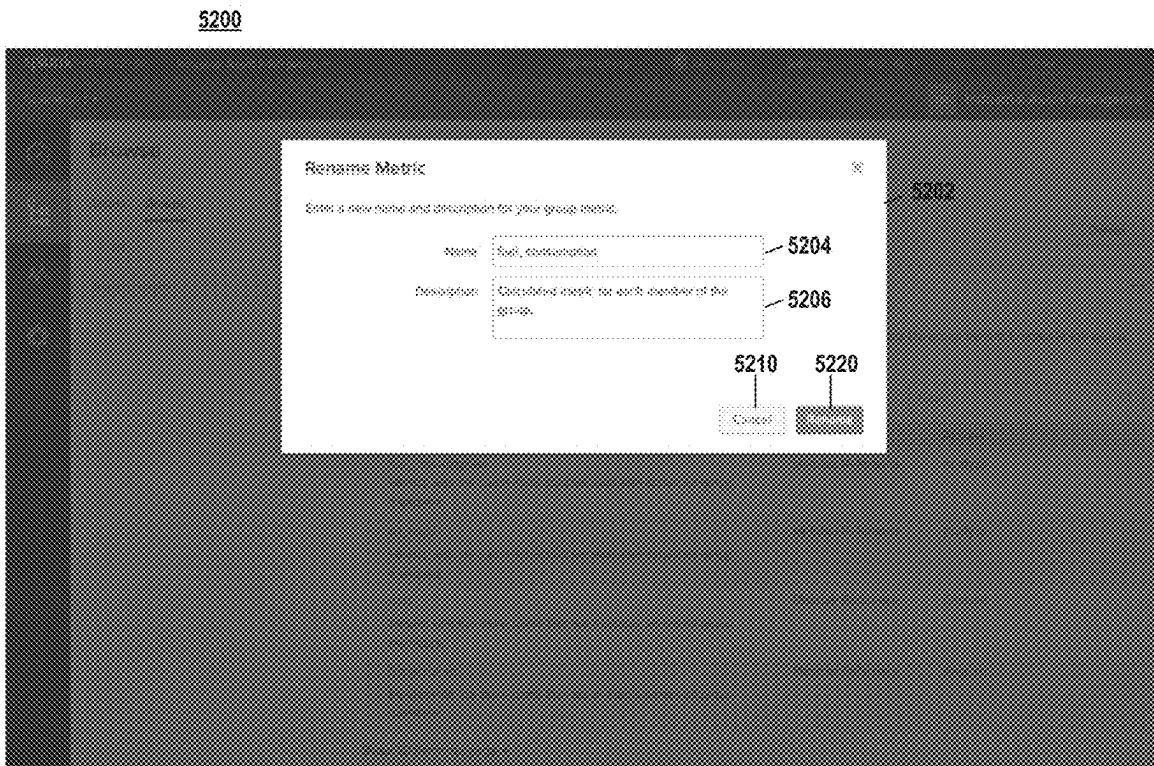


FIG. 41

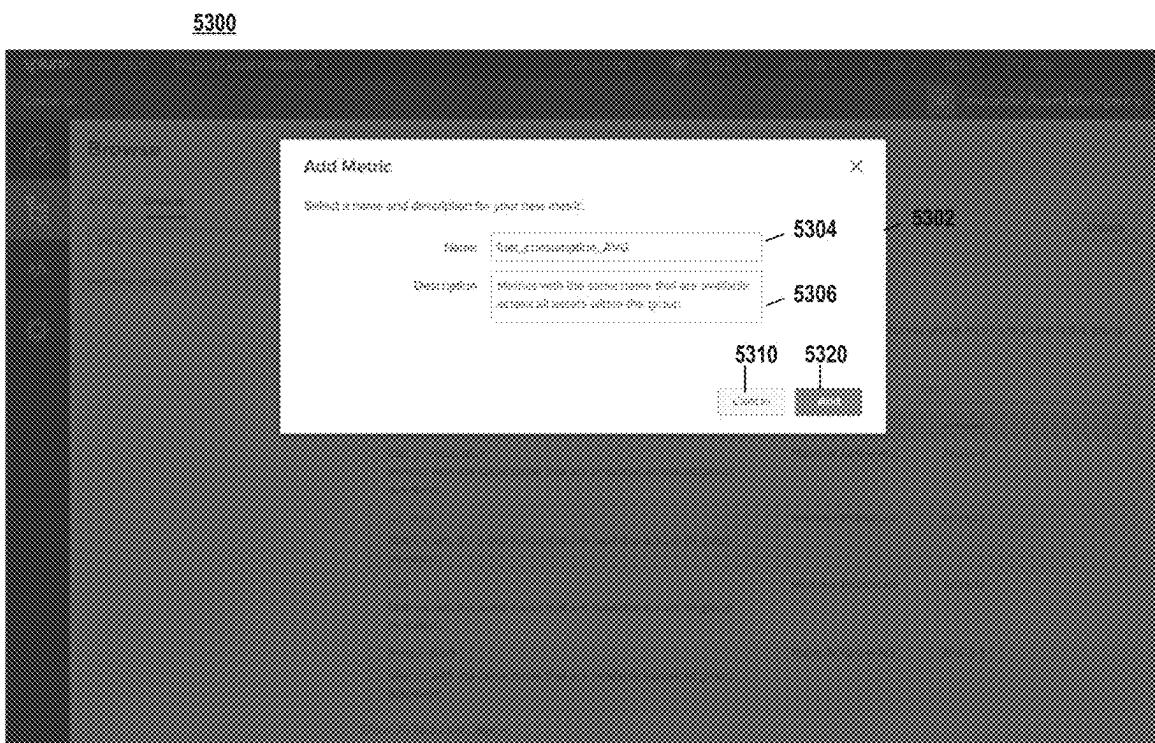


FIG. 42

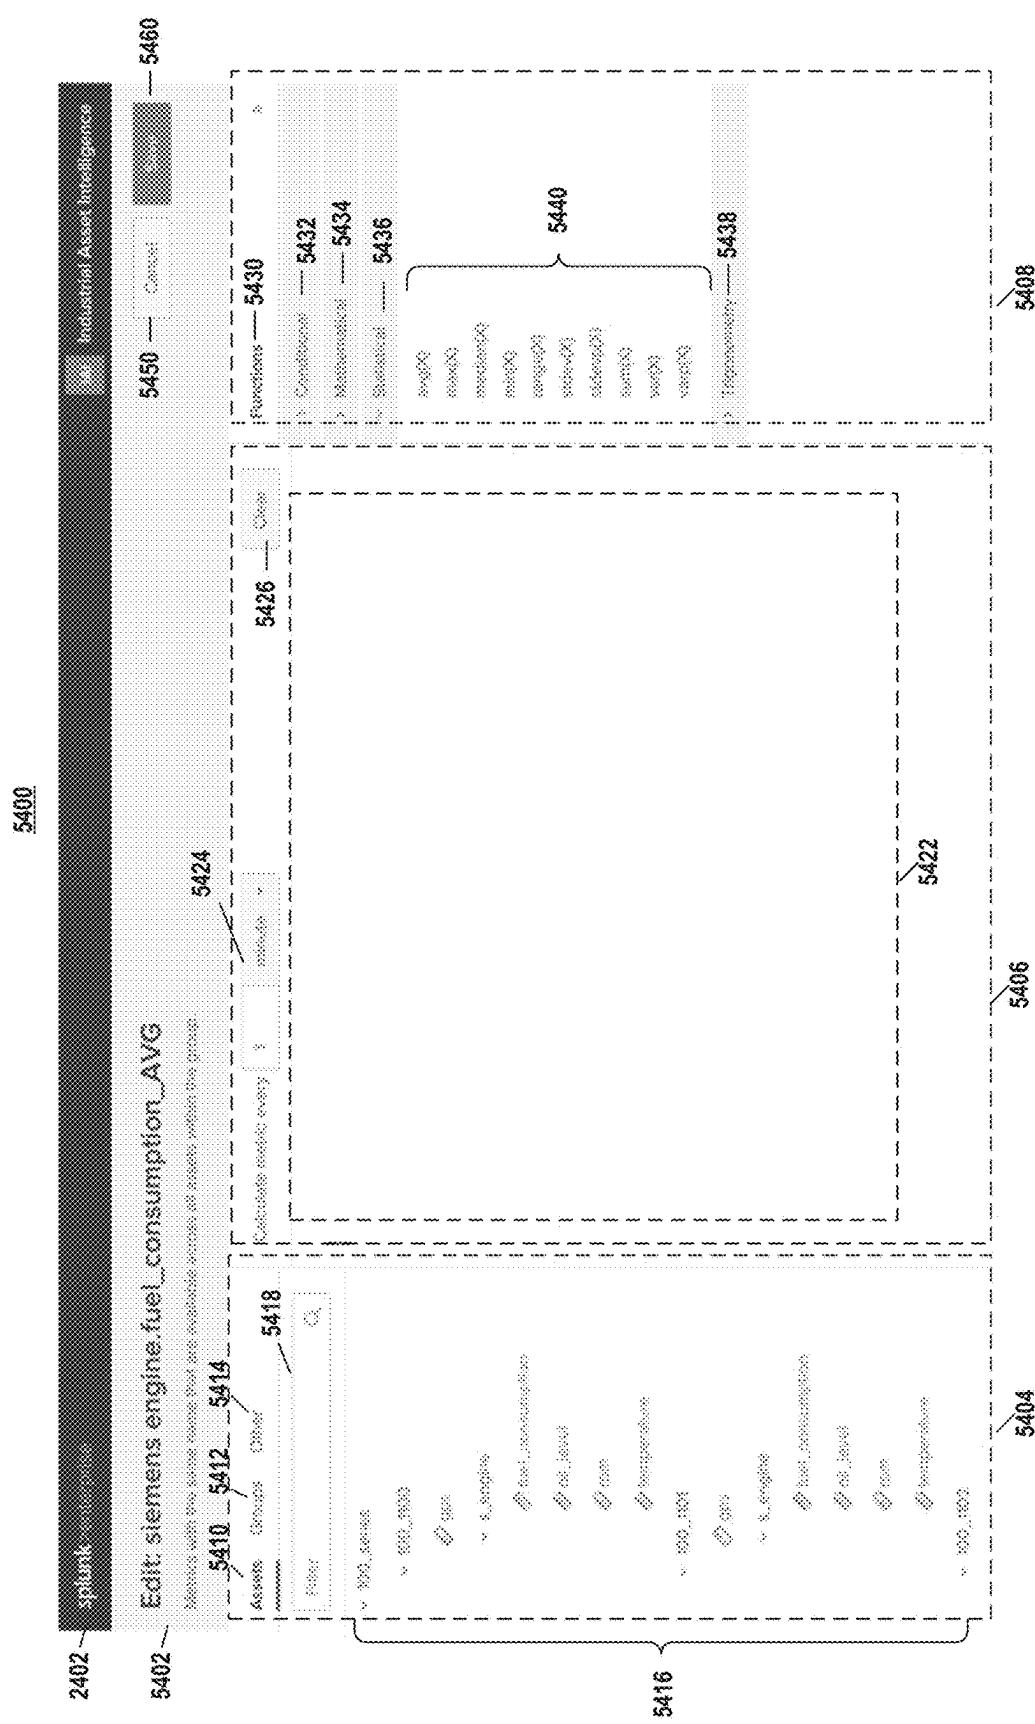


FIG. 43

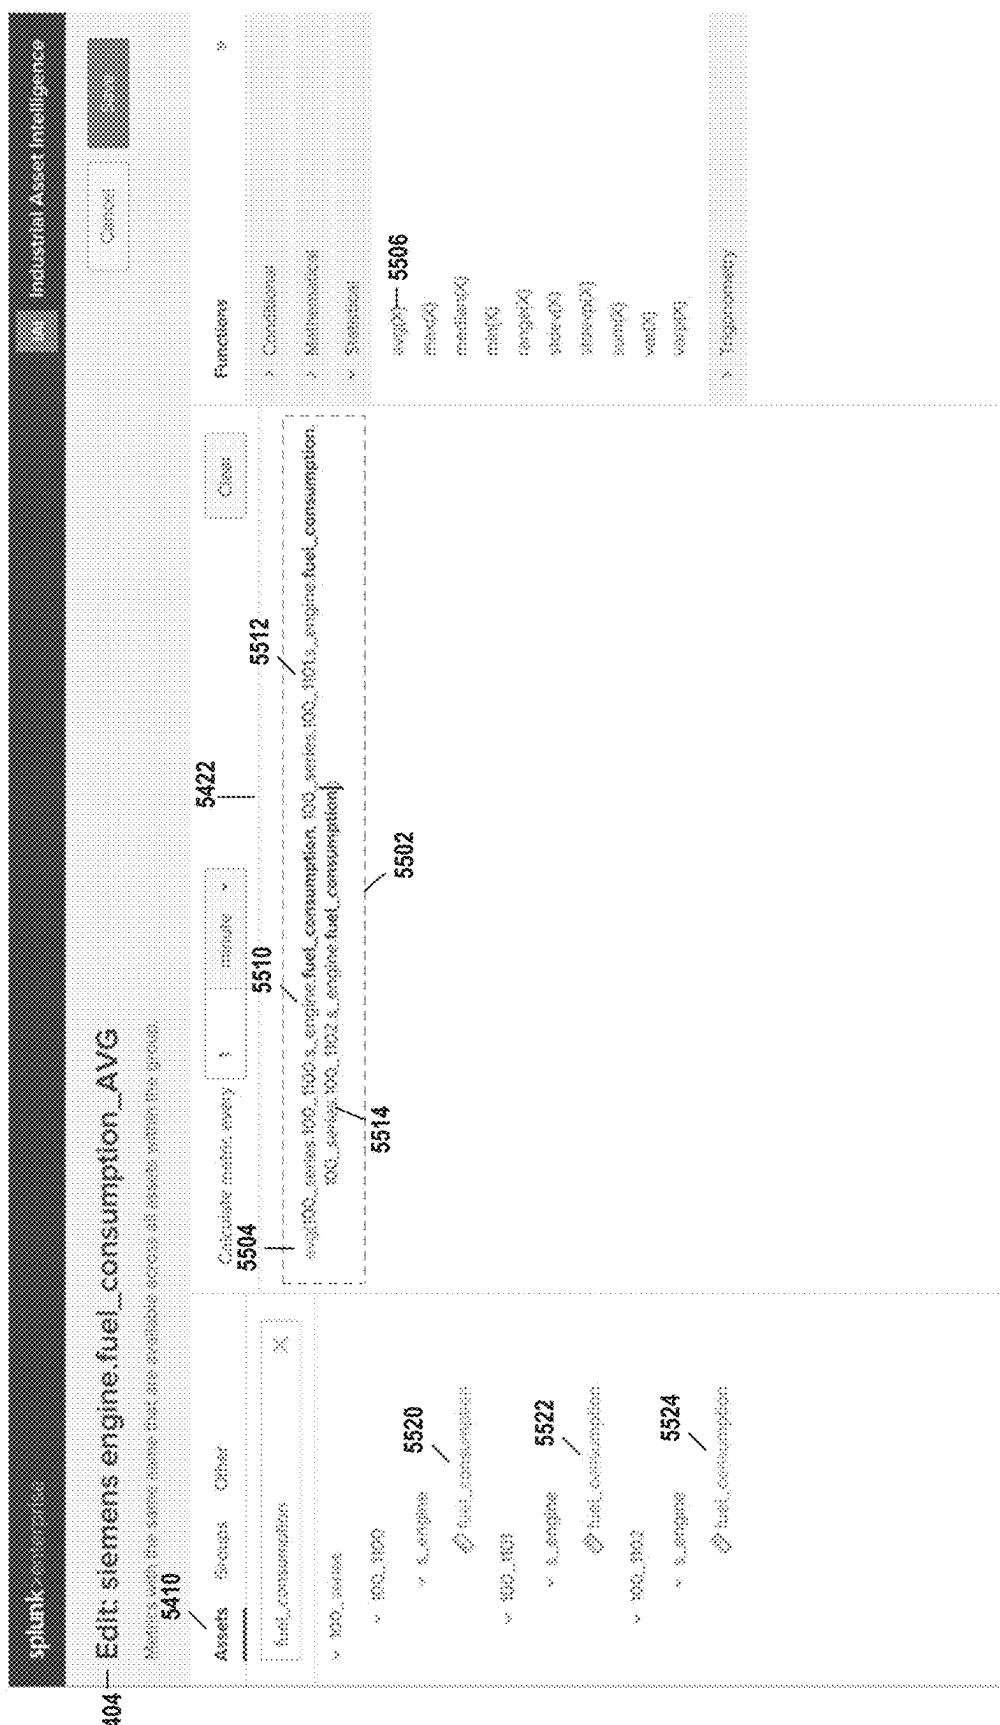


FIG. 44

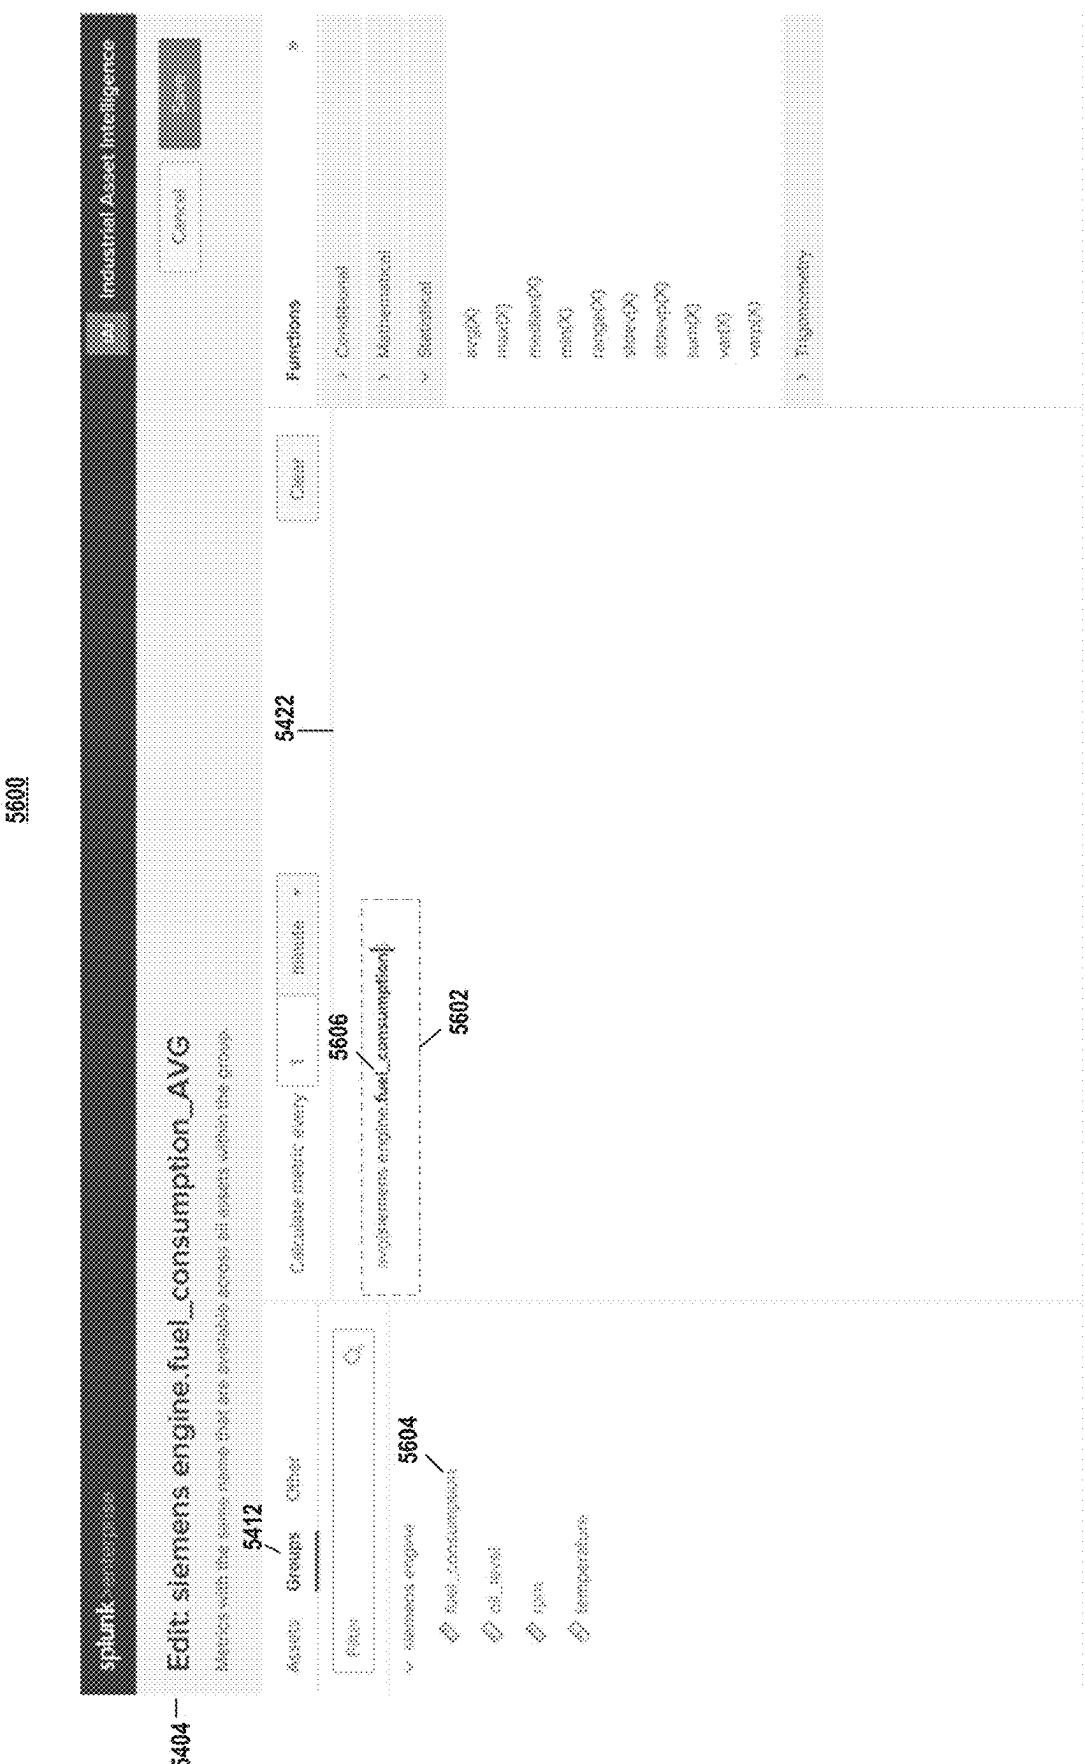
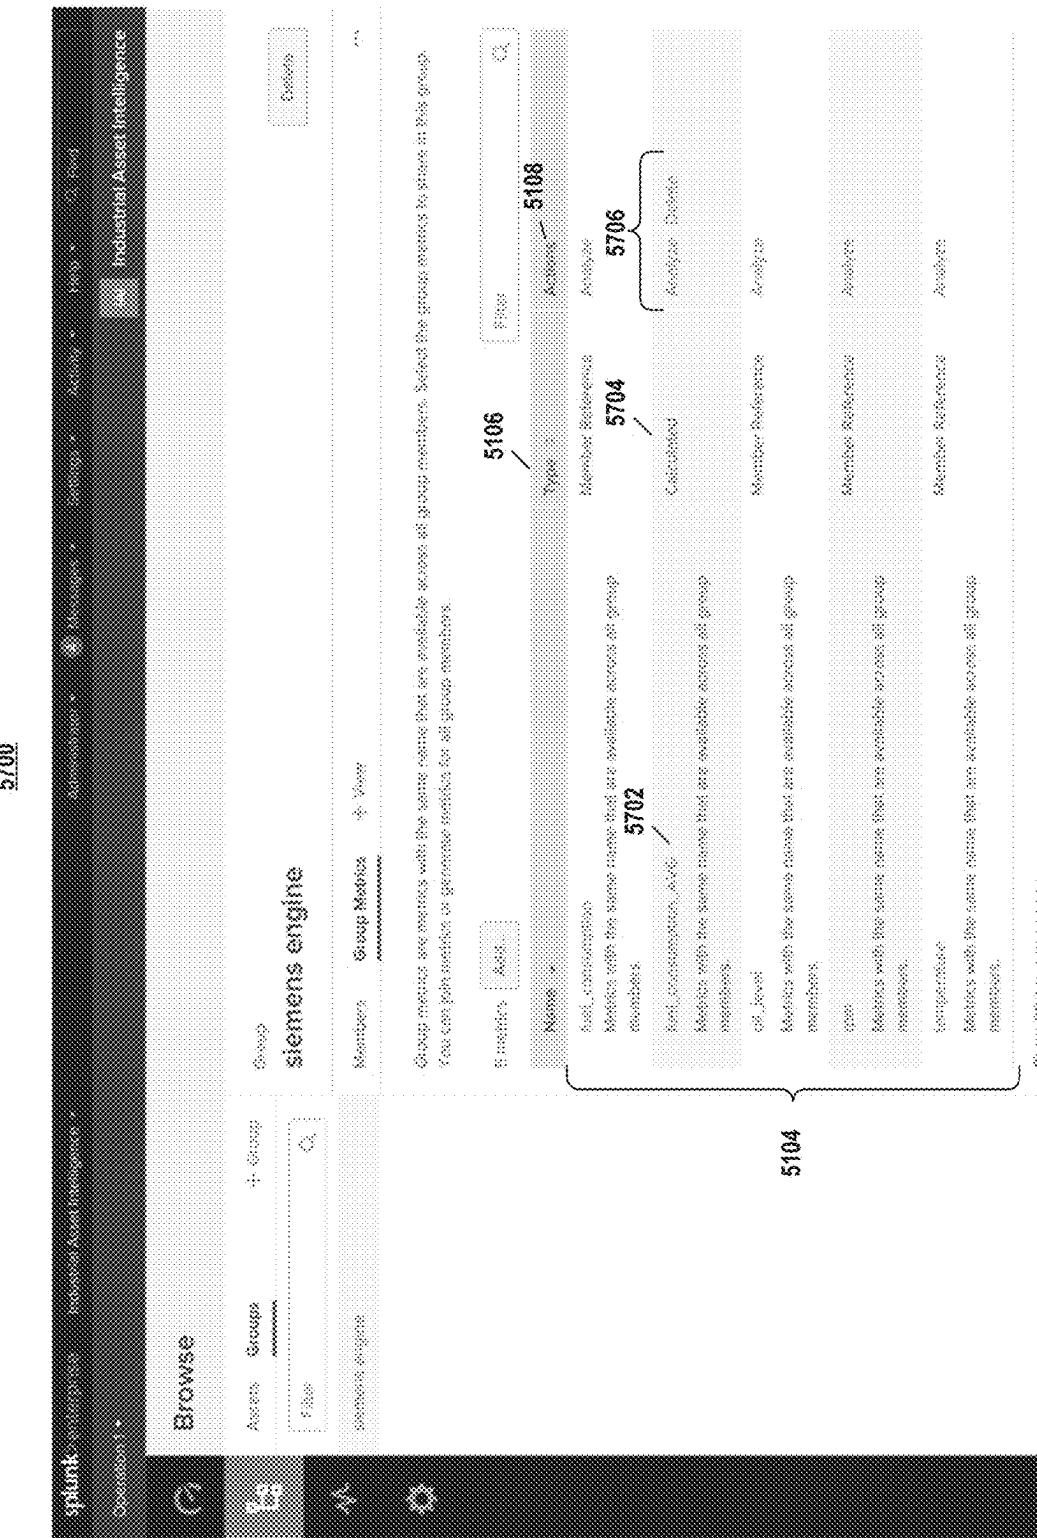


FIG. 45



EIG 46

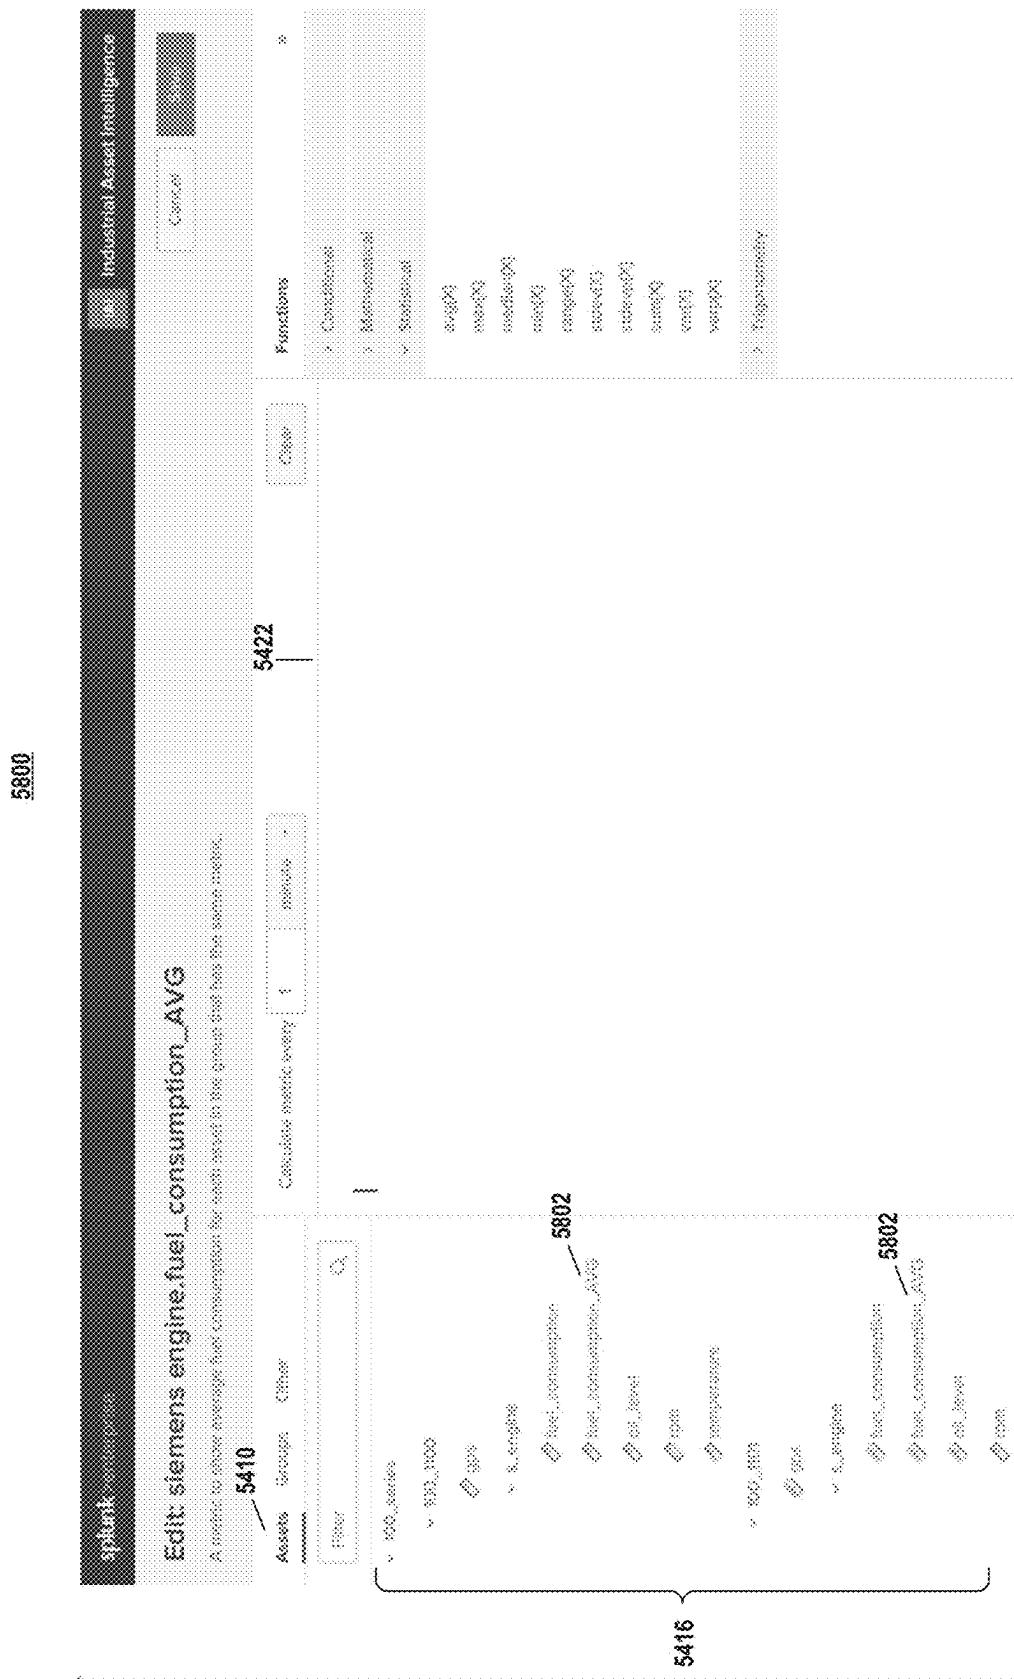


FIG. 47

5904

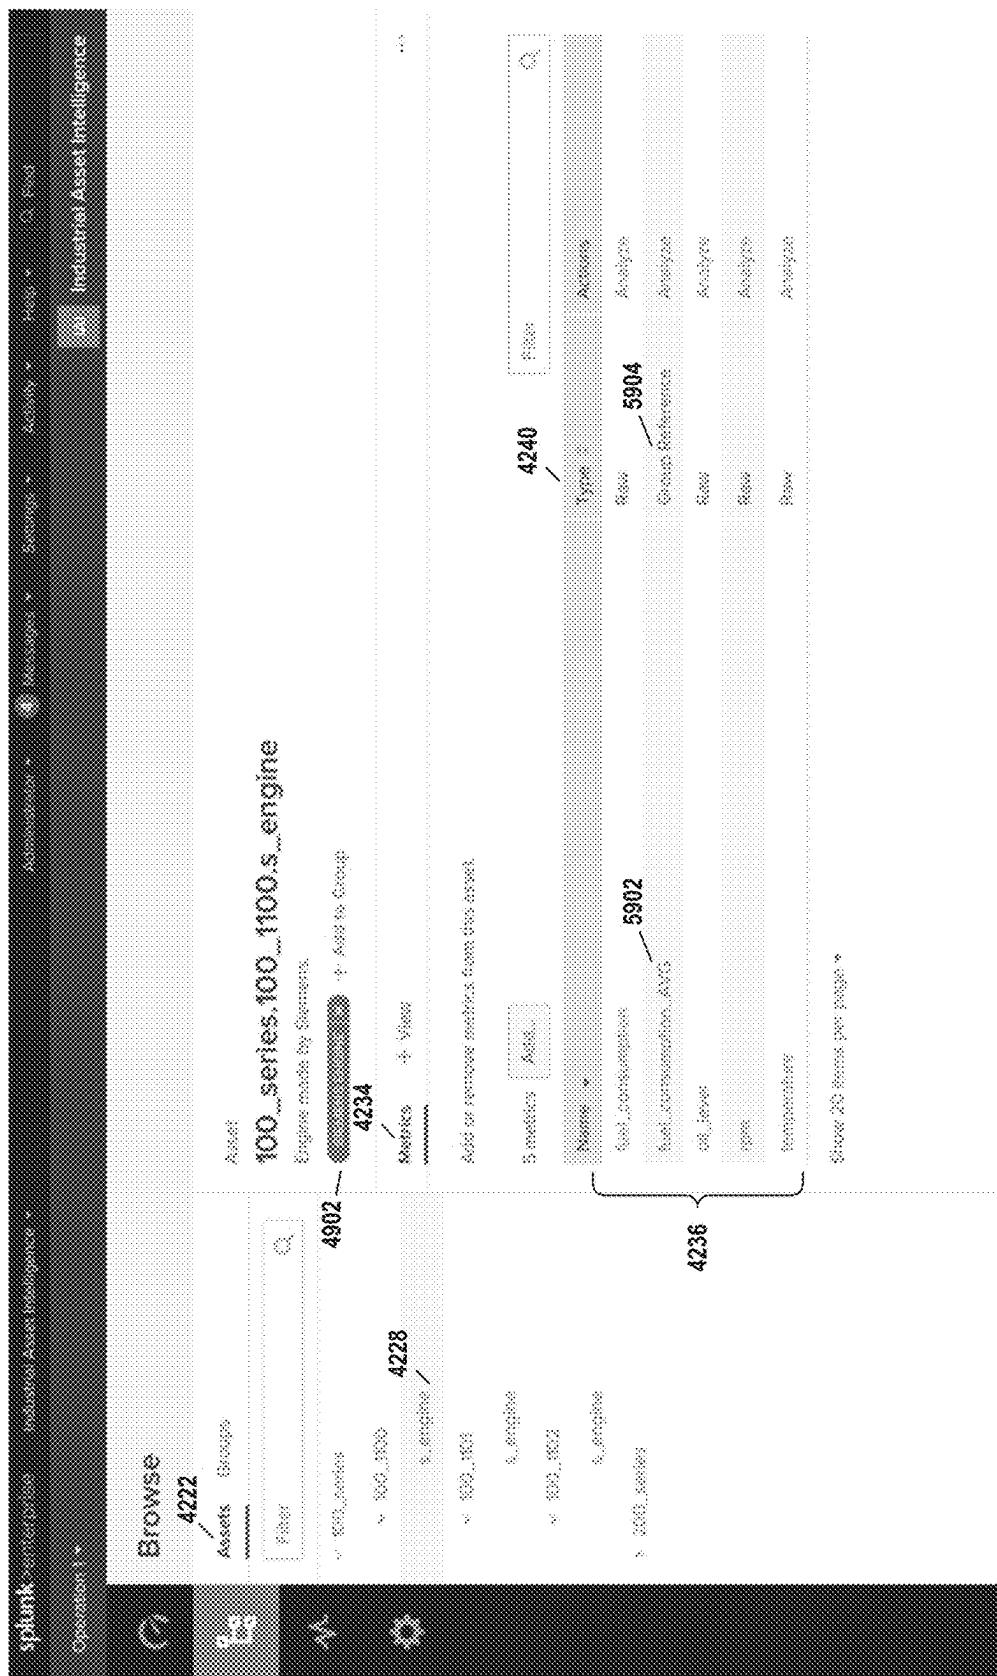


FIG. 48



FIG. 49

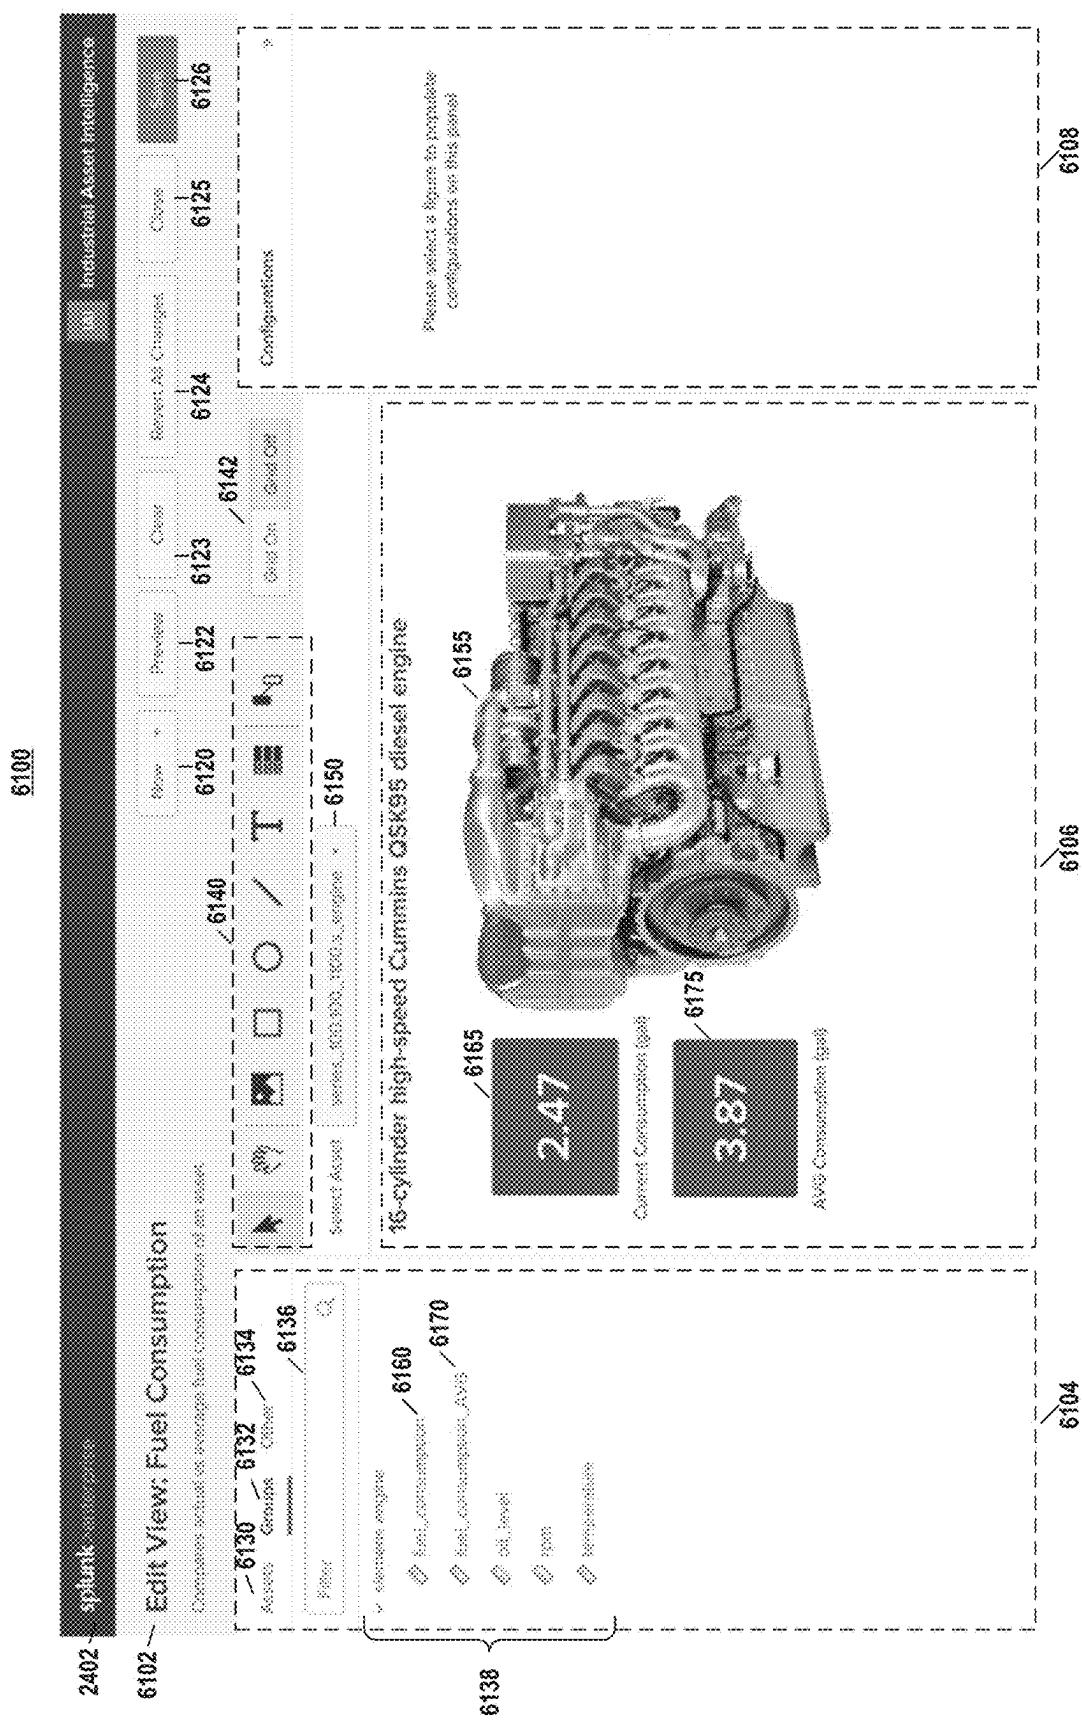


FIG. 50

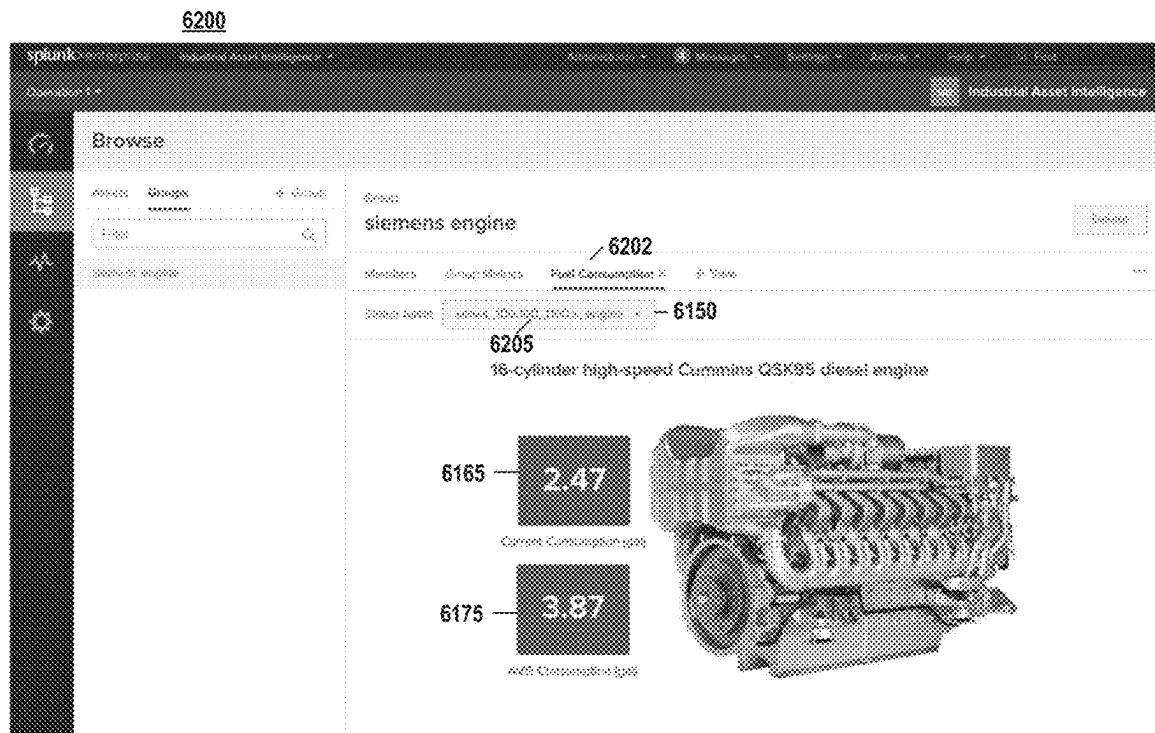


FIG. 51A

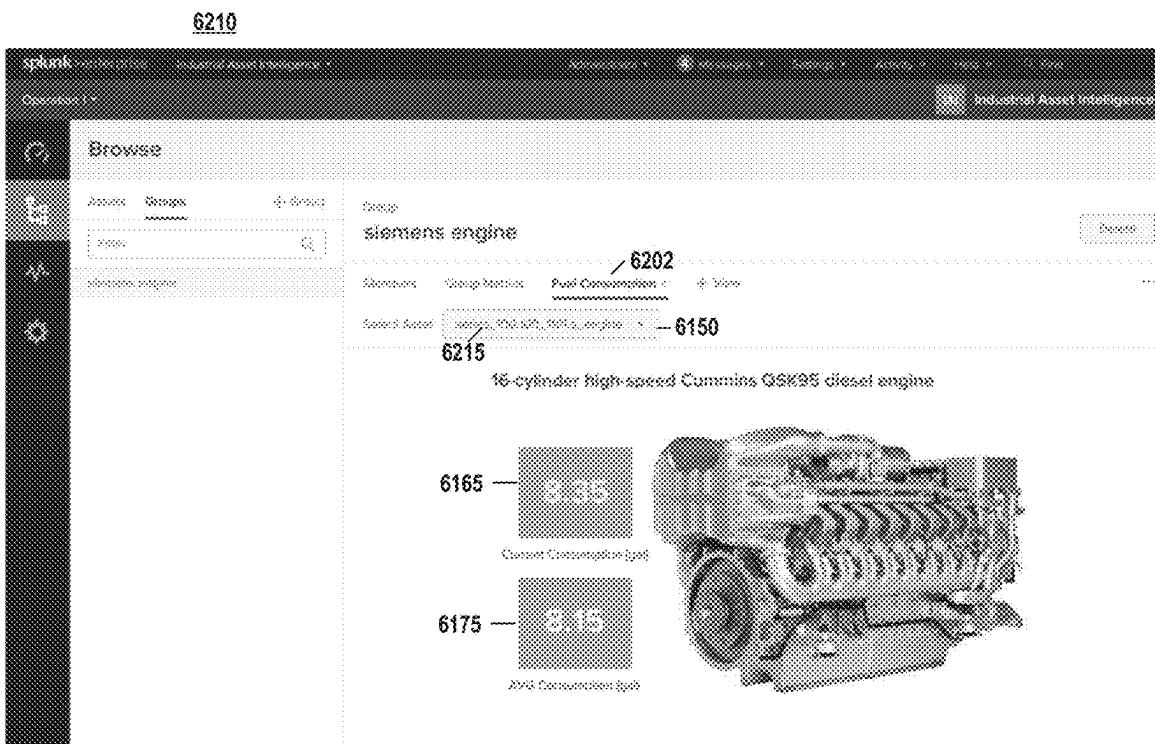


FIG. 51B

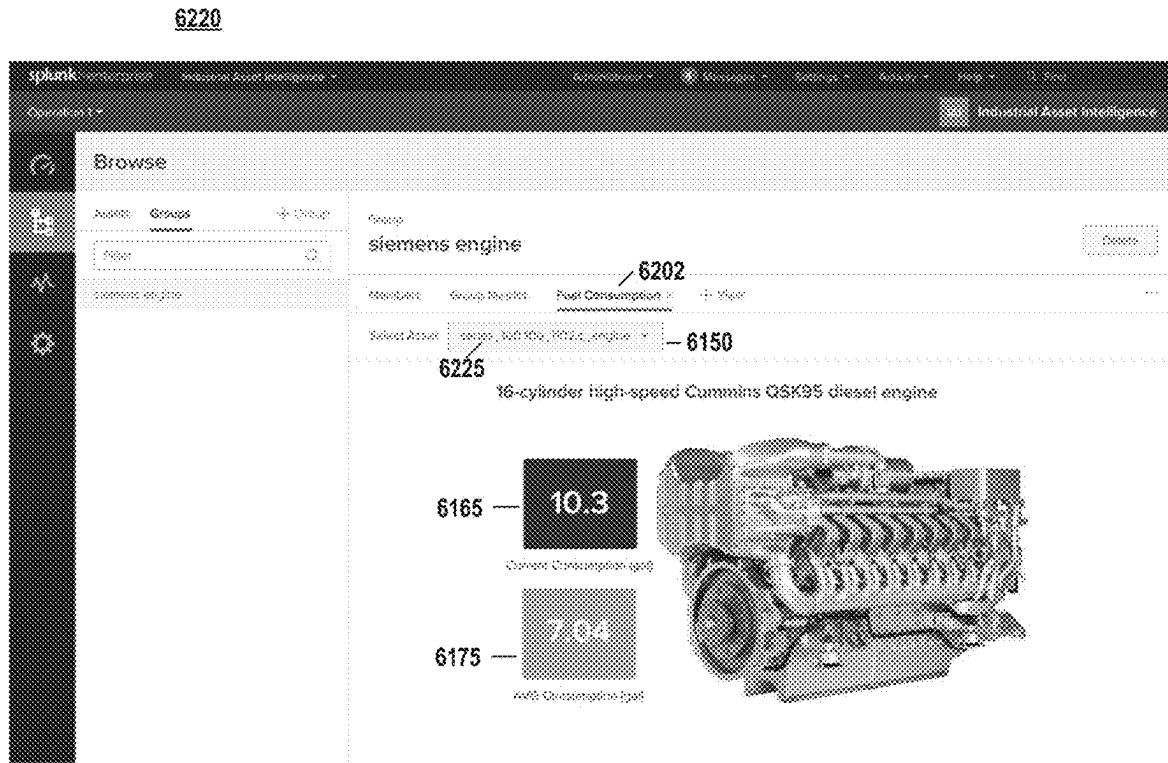


FIG. 51C

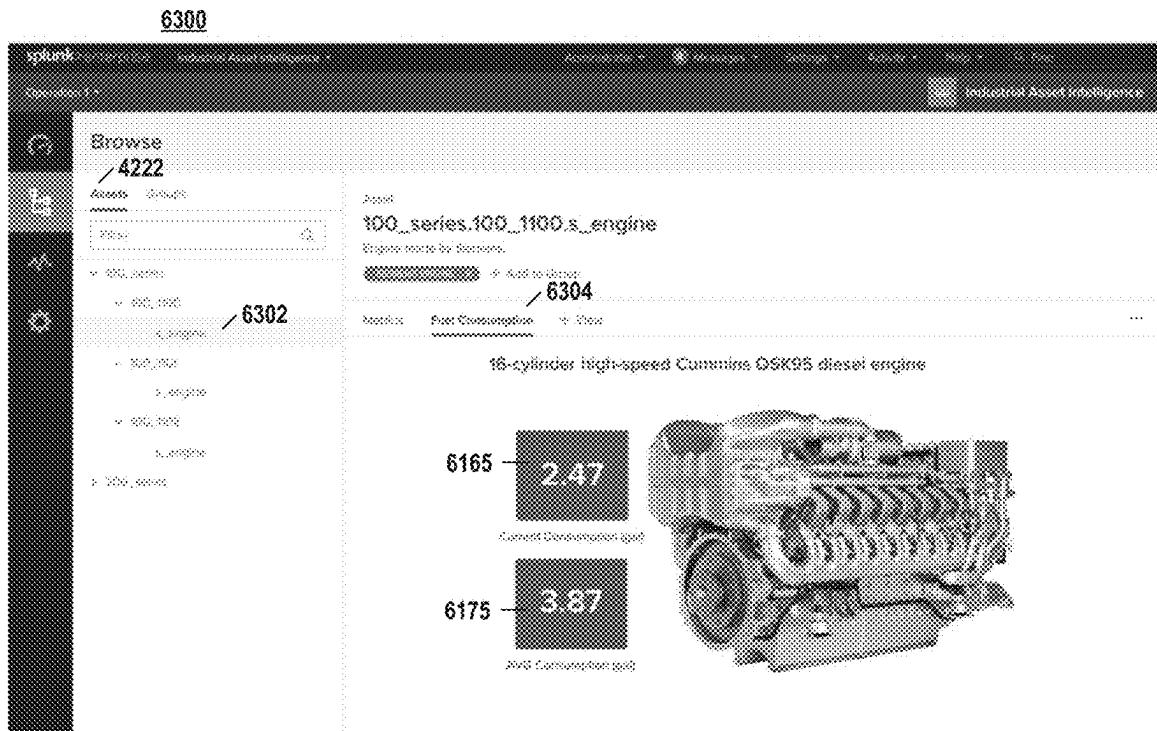


FIG. 52A

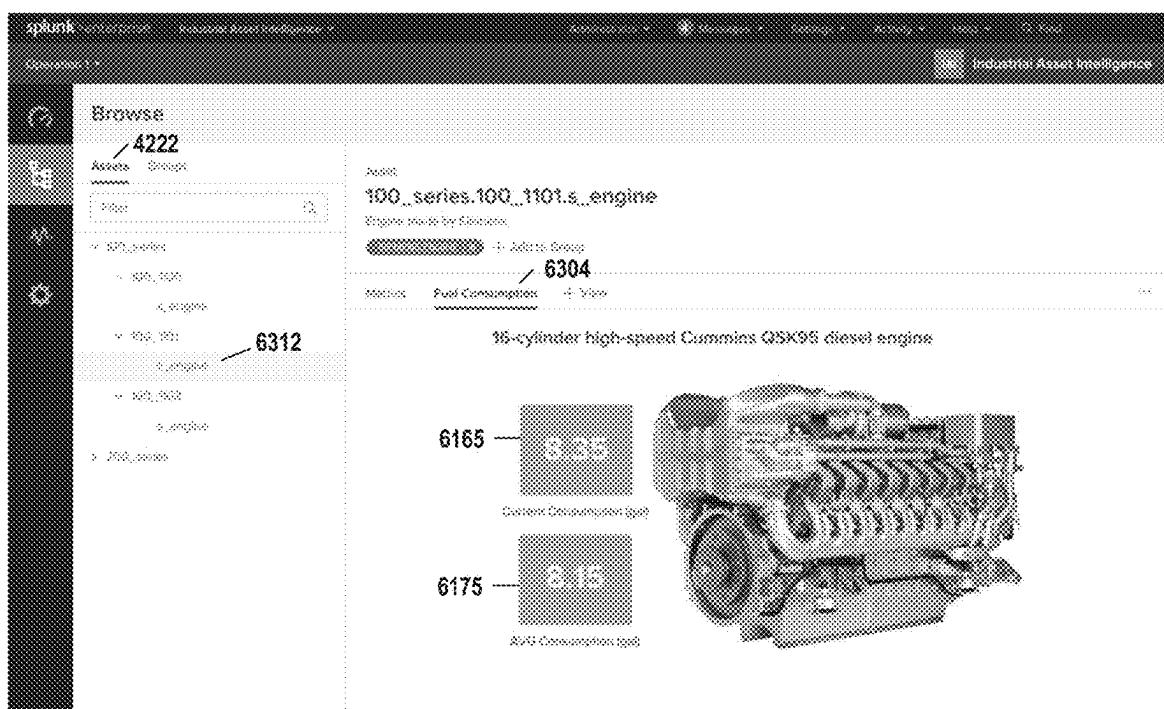


FIG. 52B

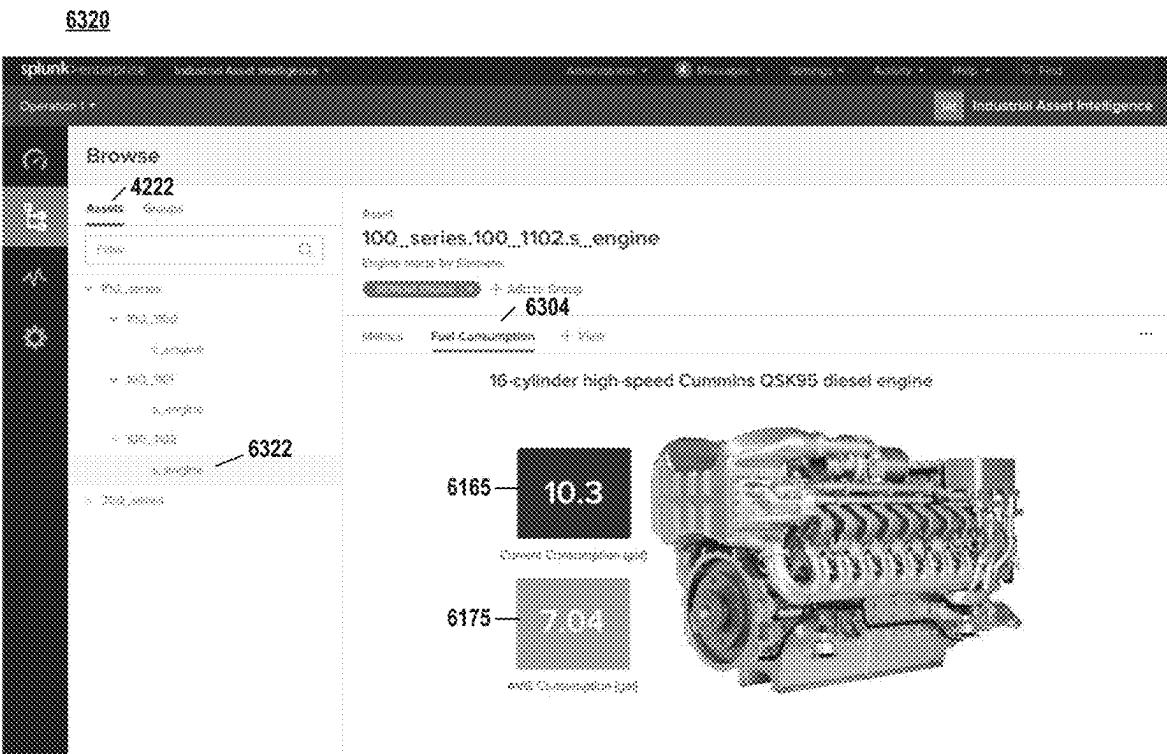


FIG. 52C

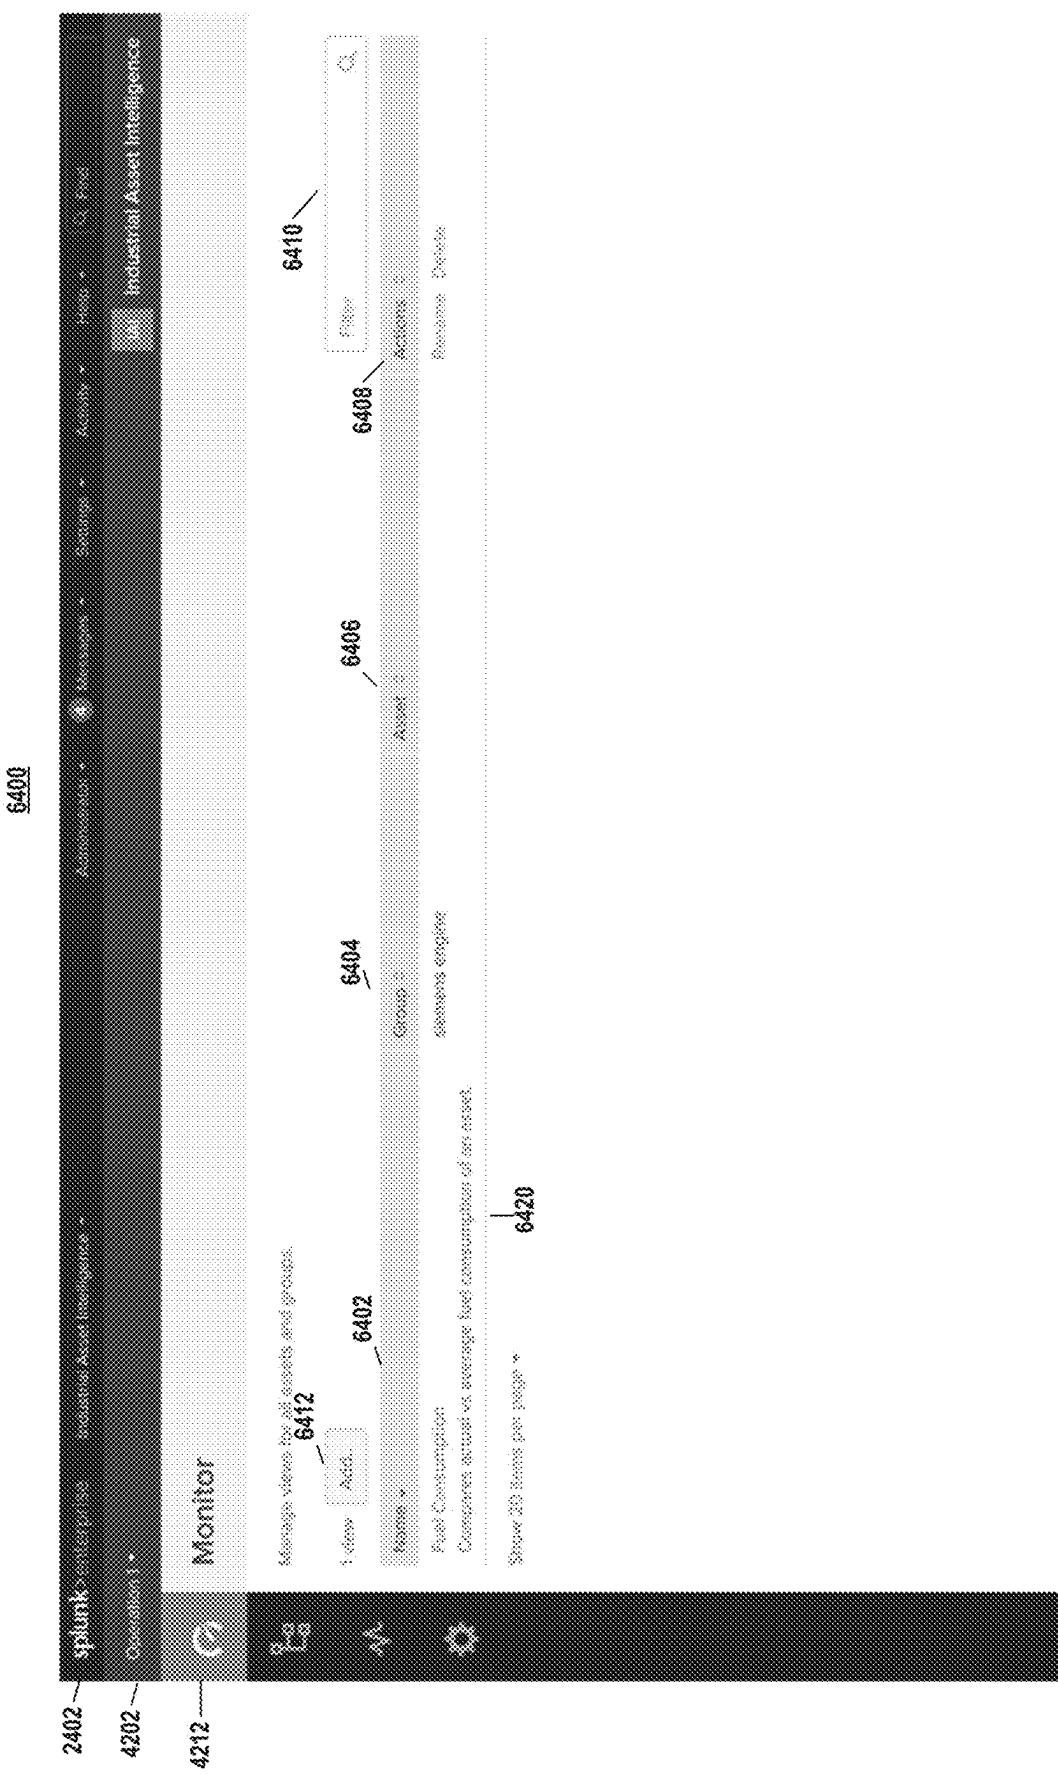


FIG. 53

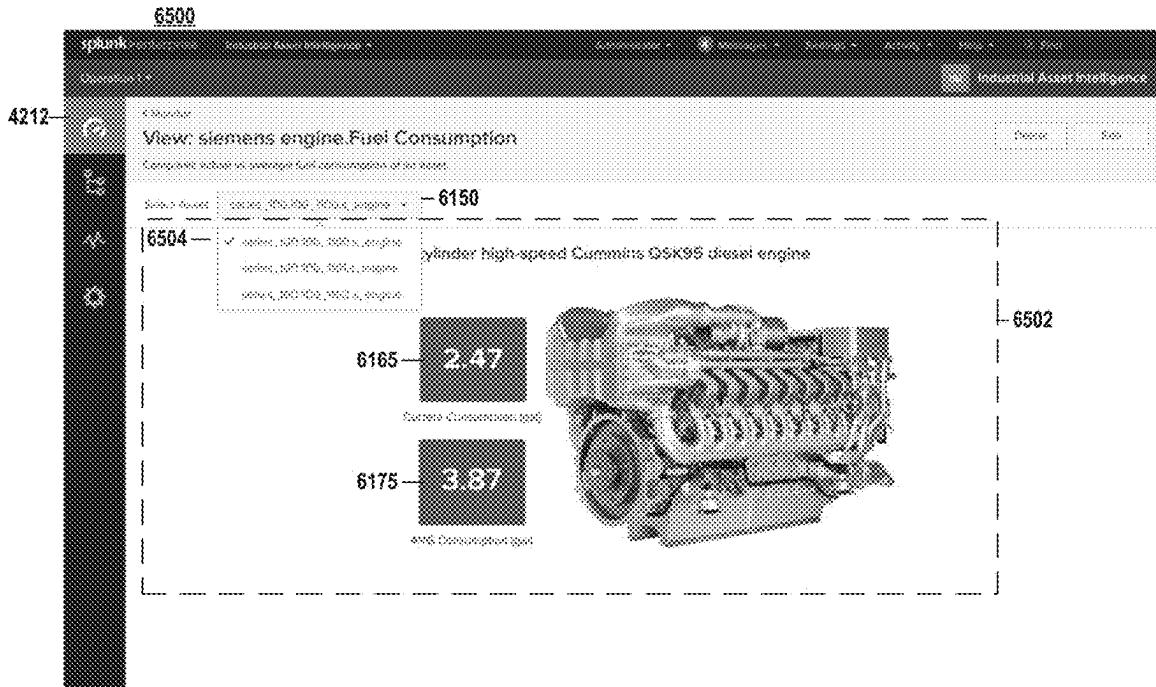


FIG. 54A

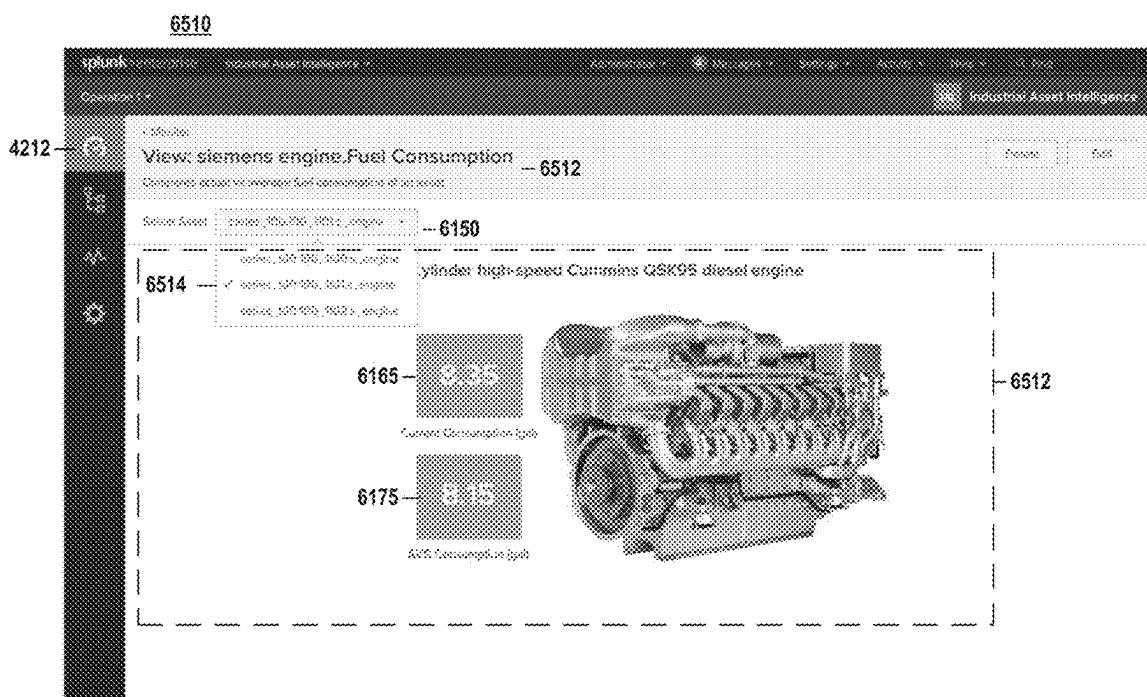


FIG. 54B

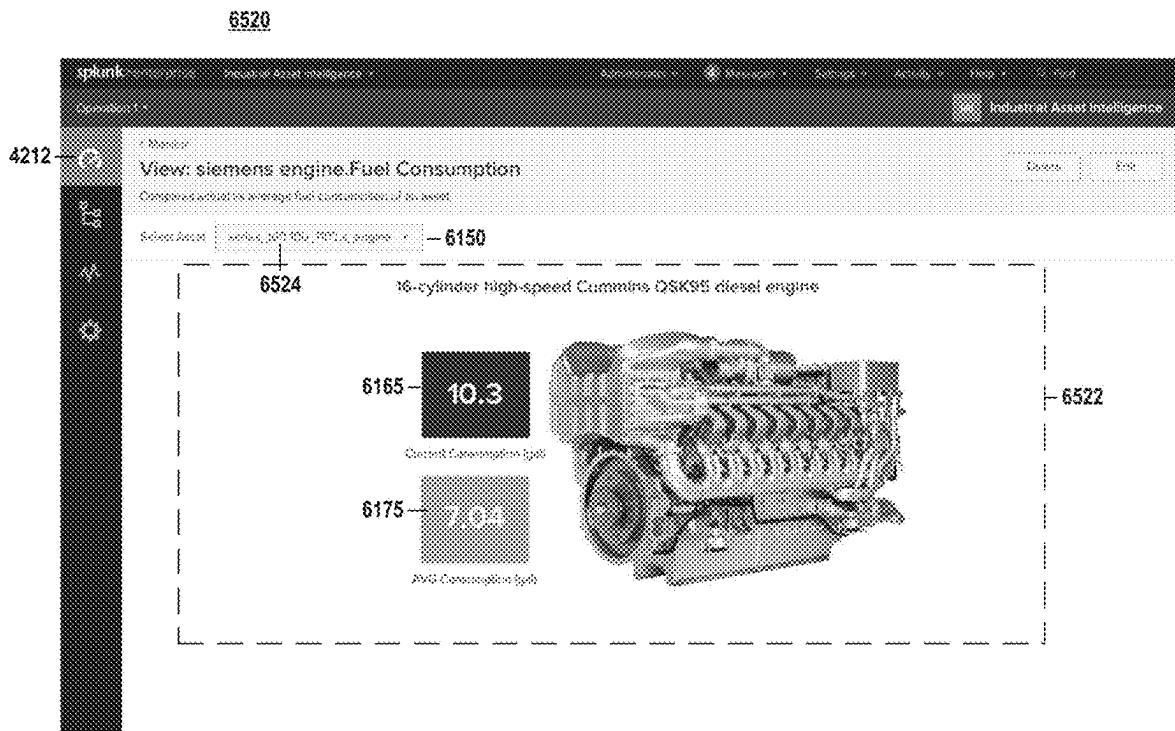


FIG. 54C

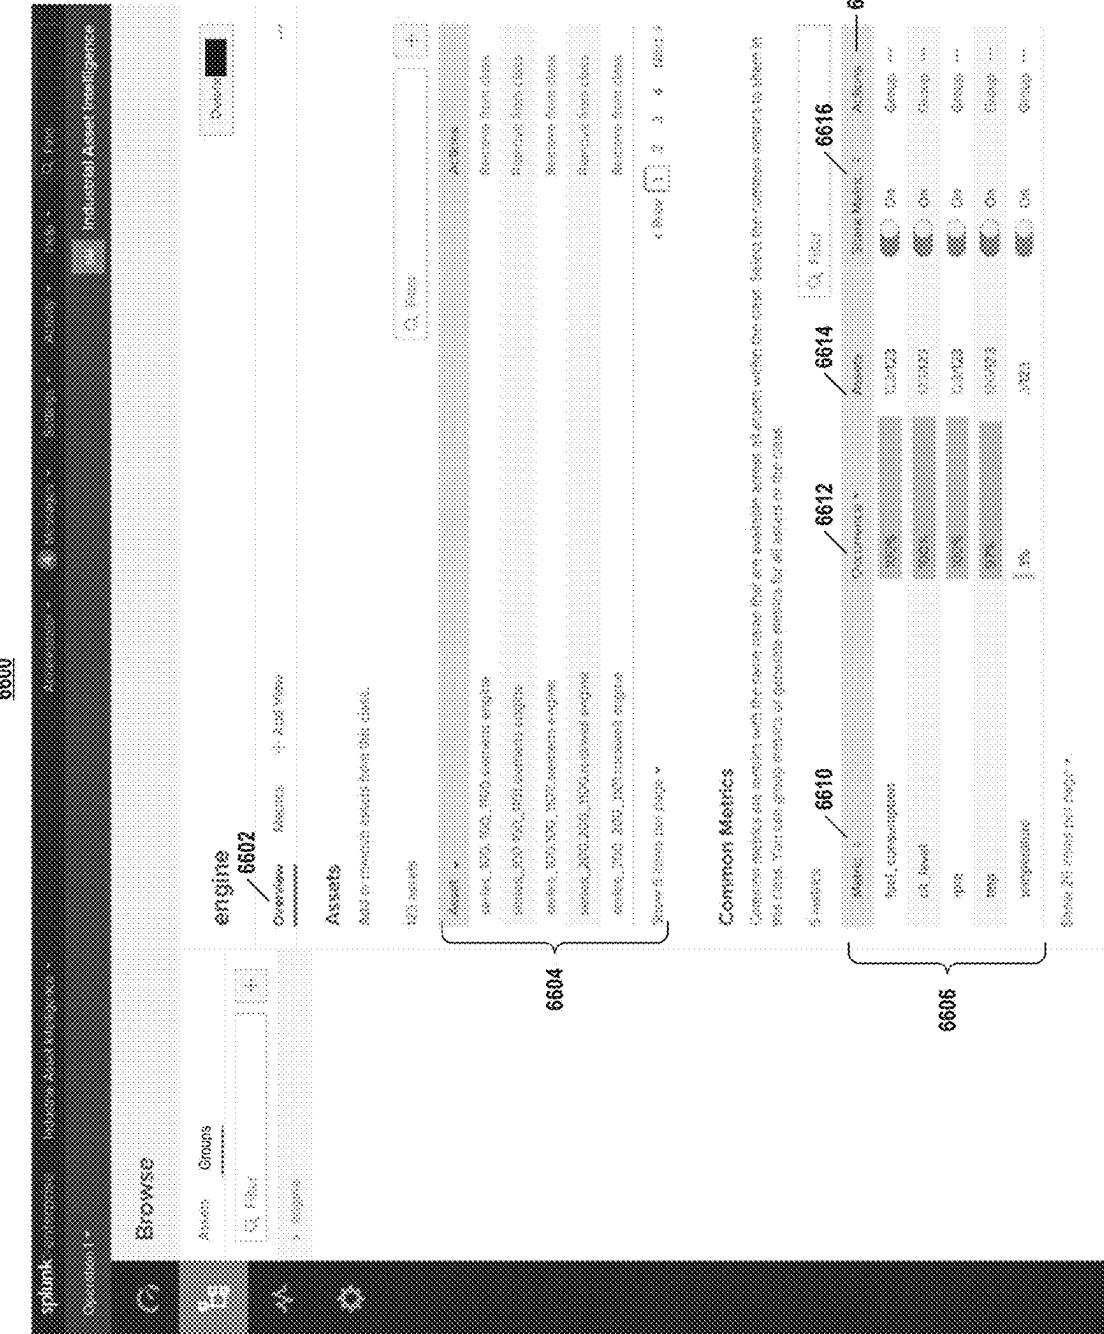
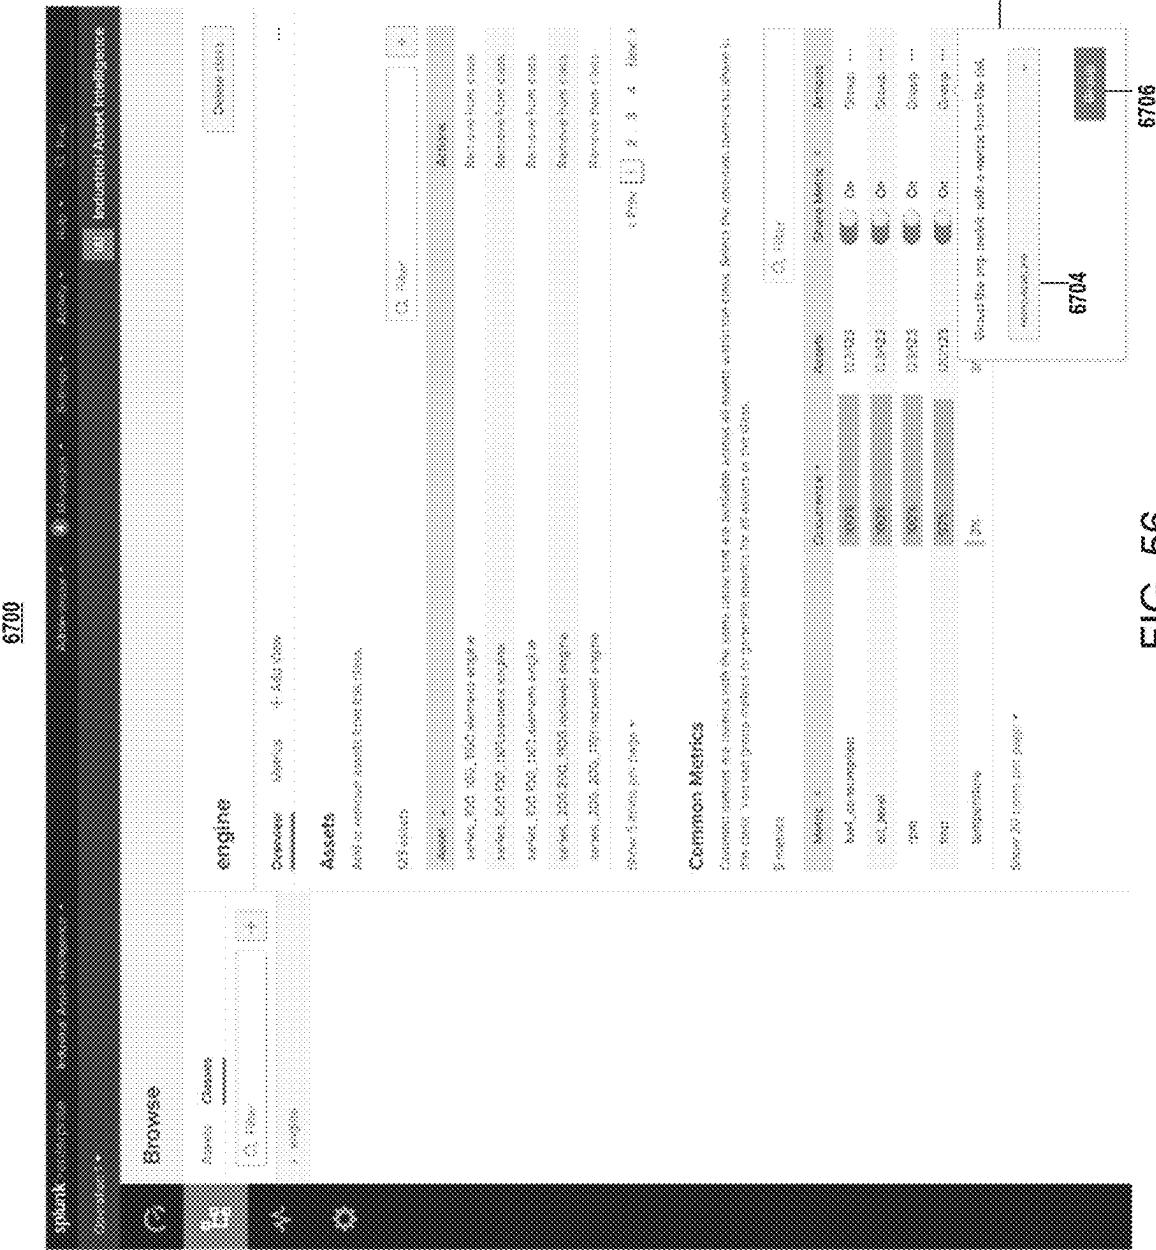


FIG. 55



56
G
II

6800

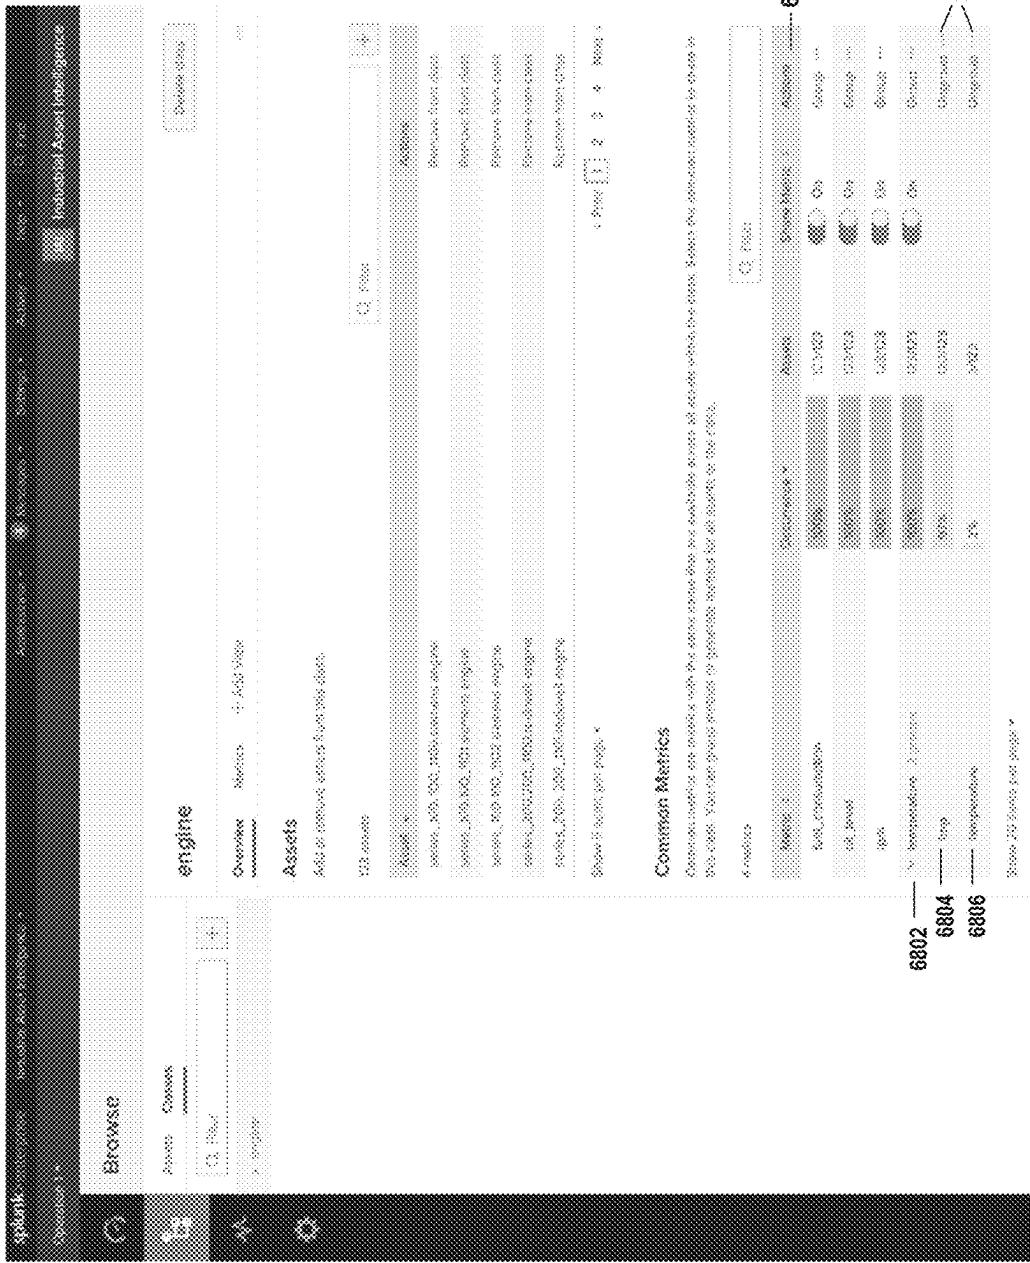


FIG. 57

ASSET GROUP MANAGEMENT USER INTERFACE

RELATED APPLICATIONS

[0001] This application is a continuation of U.S. application Ser. No. 18/206,286, filed on Jun. 6, 2023, which itself is a continuation of U.S. Pat. No. 11,703,826, filed Jul. 13, 2021, which is a continuation of U.S. Pat. No. 11,086,289, filed Jan. 15, 2020, which is a continuation of U.S. Pat. No. 10,564,622, filed Sep. 28, 2018, which is a continuation-in-part of U.S. Pat. No. 10,503,784, filed Jul. 31, 2016, which is a continuation-in-part of U.S. application Ser. No. 15/224,641, filed Jul. 31, 2016. All the above-referenced applications are incorporated by reference herein in their respective entireties.

TECHNICAL FIELD

[0002] The disclosure relates to automatic system monitoring apparatus, more particularly, to a control interface for metric definition specification for assets and asset groups driven by a search-derived asset tree hierarchy.

BACKGROUND

[0003] Modern operational systems often comprise large numbers of assets about which machine data is generated by the assets themselves or other systems and components that generate information about the asset. As the number, complexity, and sophistication of such asset-based systems, and the volume of machine data generated by and about them, increases, processing large volumes of machine-generated and machine-mediated data in an intelligent manner and effectively presenting the results of such processing continues to be a priority.

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 illustrates a networked computer environment in which an embodiment may be implemented;

[0005] FIG. 2 illustrates a block diagram of an example data intake and query system in which an embodiment may be implemented;

[0006] FIG. 3 is a flow diagram that illustrates how indexers process, index, and store data received from forwarders in accordance with the disclosed embodiments;

[0007] FIG. 4 is a flow diagram that illustrates how a search head and indexers perform a search query in accordance with the disclosed embodiments;

[0008] FIG. 5 illustrates a scenario where a common customer ID is found among log data received from three disparate sources in accordance with the disclosed embodiments;

[0009] FIG. 6A illustrates a search screen in accordance with the disclosed embodiments;

[0010] FIG. 6B illustrates a data summary dialog that enables a user to select various data sources in accordance with the disclosed embodiments;

[0011] FIG. 7 illustrates an example search query received from a client and executed by search peers in accordance with the disclosed embodiments;

[0012] FIG. 8 illustrates a block diagram of an example cloud-based data intake and query system in which an embodiment may be implemented;

[0013] FIG. 9 illustrates a block diagram of an example data intake and query system that performs searches across external data systems in accordance with the disclosed embodiments;

[0014] FIG. 10 illustrates an asset hierarchy monitoring and reporting system deployment in one embodiment.

[0015] FIGS. 11A and 11B illustrate an illustrative asset hierarchy structure.

[0016] FIG. 12 illustrates methods of an asset hierarchy monitoring and reporting system in one embodiment.

[0017] FIG. 13 illustrates a user interface display of a console function for specifying data inputs.

[0018] FIG. 14 illustrates a method for constructing an asset tree representation in control storage.

[0019] FIG. 15 illustrates a user interface display for an asset search console function.

[0020] FIG. 16 illustrates a user interface display for an asset information classification console function.

[0021] FIG. 17 illustrates a user interface for an asset tree display console function.

[0022] FIG. 18 illustrates a user interface display for a metrics console function.

[0023] FIG. 19 illustrates a user interface display for a metrics configuration console function.

[0024] FIG. 20 illustrates a user interface display for a metrics condition and alerts console function.

[0025] FIG. 21 illustrates a user interface display for an actions console function.

[0026] FIG. 22 illustrates a user interface display for creating or editing a custom monitoring or reporting presentation for an asset tree.

[0027] FIG. 23 illustrates a user interface display of a custom asset tree presentation.

[0028] FIG. 24 illustrates a user interface display for a metrics view.

[0029] FIG. 25 illustrates a user interface display for a conditions and alerts view.

[0030] FIG. 26 illustrates a user interface display for a diagnostics view.

[0031] FIG. 27 illustrates a user interface display for a map view of asset tree data.

[0032] FIG. 28 illustrates a user interface display for a map view of asset tree data with a timeline.

[0033] FIG. 29 illustrates methods of an asset hierarchy monitoring and reporting system in one embodiment.

[0034] FIGS. 30A-B illustrate methods for generation and management of asset groups and

[0035] asset group metrics in view of an asset tree representation in control storage.

[0036] FIG. 31 illustrates a user interface display for an asset group console function.

[0037] FIG. 32 illustrates a user interface display for an asset group console function for assign a new asset group.

[0038] FIG. 33 illustrates a user interface display for an asset group creation console function for specifying an asset group name.

[0039] FIG. 34 illustrates a user interface display for an asset group console function for adding a new asset to an asset group.

[0040] FIG. 35 illustrates a user interface display an asset group console function for specifying assets to add to an asset group.

[0041] FIG. 36 illustrates a user interface display for an asset group members console.

- [0042] FIG. 37 illustrates a user interface display for adding an in-focus asset to an asset group.
- [0043] FIG. 38 illustrates a user interface display for addition of an asset member to an asset group.
- [0044] FIG. 39 illustrates a user interface display for an asset group browsing console for addition of assets to an asset group.
- [0045] FIG. 40 illustrates a user interface display for group metrics of an asset group console.
- [0046] FIG. 41 illustrates a user interface display for renaming a group metric of an asset group.
- [0047] FIG. 42 illustrates a user interface display for adding a group metric to an asset group.
- [0048] FIG. 43 illustrates a user interface display for configuring a group metric of an asset group console.
- [0049] FIG. 44 illustrates a user interface display for configuring a group metric of an asset group console with editable group metric definitions based on asset metrics.
- [0050] FIG. 45 illustrates a user interface display for configuring a group metric of an asset group console with editable group metric definitions based on asset group metrics.
- [0051] FIG. 46 illustrates a user interface display for an asset group console with group metrics browsing.
- [0052] FIG. 47 illustrates a user interface display for an asset group metric configuration console.
- [0053] FIG. 48 illustrates a user interface display for an asset group console with group metrics replicated to an asset.
- [0054] FIG. 49 illustrates a user interface display for adding an asset group view for an asset group.
- [0055] FIG. 50 illustrates a user interface display for configuring an asset group view via an asset group view console.
- [0056] FIGS. 51A-51C illustrate user interface displays for asset group view browsing at the asset group level.
- [0057] FIGS. 52A-52C illustrate user interface displays for asset group view browsing at the asset level.
- [0058] FIG. 53 illustrates a user interface display for asset group view monitoring configuration console.
- [0059] FIGS. 54A-54C illustrate user interface displays for asset group view monitoring.
- [0060] FIG. 55 illustrates a user interface display for an asset group management console.
- [0061] FIG. 56 illustrates a user interface display for joining group metrics of an asset group.
- [0062] FIG. 57 illustrates a user interface display for joint group metrics of an asset group.

DETAILED DESCRIPTION

[0063] Modern data centers and other computing and automation environments can comprise anywhere from a few host computer systems to thousands of systems configured to process data, service requests from remote clients, and perform numerous other computational tasks. During operation, various components within these computing environments often generate significant volumes of machine-generated data. For example, machine data is generated by various components in the information technology (IT) environments, such as servers, sensors, routers, mobile devices, Internet of Things (IoT) devices, etc. Machine-generated data can include system logs, network packet data, sensor data, application program data, error logs, stack traces, system performance data, etc. In general, machine-

generated data can also include performance data, diagnostic information, and many other types of data that can be analyzed to diagnose performance problems, monitor user interactions, and to derive other insights.

[0064] A number of tools are available to analyze machine data, that is, machine-generated data. In order to reduce the size of the potentially vast amount of machine data that may be generated, many of these tools typically pre-process the data based on anticipated data-analysis needs. For example, pre-specified data items may be extracted from the machine data and stored in a database to facilitate efficient retrieval and analysis of those data items at search time. However, the rest of the machine data typically is not saved and discarded during pre-processing. As storage capacity becomes progressively cheaper and more plentiful, there are fewer incentives to discard these portions of machine data and many reasons to retain more of the data.

[0065] This plentiful storage capacity is presently making it feasible to store massive quantities of minimally processed machine data for later retrieval and analysis. In general, storing minimally processed machine data and performing analysis operations at search time can provide greater flexibility because it enables an analyst to search all of the machine data, instead of searching only a pre-specified set of data items. This may enable an analyst to investigate different aspects of the machine data that previously were unavailable for analysis.

[0066] However, analyzing and searching massive quantities of machine data presents a number of challenges. For example, a data center, servers, or network appliances may generate many different types and formats of machine data (e.g., system logs, network packet data (e.g., wire data, etc.), sensor data, application program data, error logs, stack traces, system performance data, operating system data, virtualization data, etc.) from thousands of different components, which can collectively be very time-consuming to analyze. In another example, mobile devices may generate large amounts of information relating to data accesses, application performance, operating system performance, network performance, etc. There can be millions of mobile devices that report these types of information.

[0067] These challenges can be addressed by using an event-based data intake and query system, such as the SPLUNK® ENTERPRISE system developed by Splunk Inc. of San Francisco, California. The SPLUNK® ENTERPRISE system is the leading platform for providing real-time operational intelligence that enables organizations to collect, index, and search machine-generated data from various websites, applications, servers, networks, and mobile devices that power their businesses. The SPLUNK® ENTERPRISE system is particularly useful for analyzing data which is commonly found in system log files, network data, and other data input sources. Although many of the techniques described herein are explained with reference to a data intake and query system similar to the SPLUNK® ENTERPRISE system, these techniques are also applicable to other types of data systems.

[0068] In the SPLUNK® ENTERPRISE system, machine-generated data are collected and stored as “events”. An event comprises a portion of the machine-generated data and is associated with a specific point in time. For example, events may be derived from “time series data,” where the time series data comprises a sequence of data points (e.g., performance measurements from a computer system, etc.)

that are associated with successive points in time. In general, each event can be associated with a timestamp that is derived from the raw data in the event, determined through interpolation between temporally proximate events having known timestamps, or determined based on other configurable rules for associating timestamps with events, etc.

[0069] In some instances, machine data can have a pre-defined format, where data items with specific data formats are stored at predefined locations in the data. For example, the machine data may include data stored as fields in a database table. In other instances, machine data may not have a predefined format, that is, the data is not at fixed, predefined locations, but the data does have repeatable patterns and is not random. This means that some machine data can comprise various data items of different data types and that may be stored at different locations within the data. For example, when the data source is an operating system log, an event can include one or more lines from the operating system log containing raw data that includes different types of performance and diagnostic information associated with a specific point in time.

[0070] Examples of components which may generate machine data from which events can be derived include, but are not limited to, web servers, application servers, databases, firewalls, routers, operating systems, and software applications that execute on computer systems, mobile devices, sensors, Internet of Things (IoT) devices, etc. The data generated by such data sources can include, for example and without limitation, server log files, activity log files, configuration files, messages, network packet data, performance measurements, sensor measurements, etc.

[0071] The SPLUNK® ENTERPRISE system uses flexible schema to specify how to extract information from the event data. A flexible schema may be developed and redefined as needed. Note that a flexible schema may be applied to event data “on the fly,” when it is needed (e.g., at search time, index time, ingestion time, etc.). When the schema is not applied to event data until search time it may be referred to as a “late-binding schema.”

[0072] During operation, the SPLUNK® ENTERPRISE system starts with raw input data (e.g., one or more system logs, streams of network packet data, sensor data, application program data, error logs, stack traces, system performance data, etc.) The system divides this raw data into blocks (e.g., buckets of data, each associated with a specific time frame, etc.), and parses the raw data to produce timestamped events. The system stores the timestamped events in a data store. The system enables users to run queries against the stored data to, for example, retrieve events that meet criteria specified in a query, such as containing certain keywords or having specific values in defined fields. As used herein throughout, data that is part of an event is referred to as “event data”. In this context, the term “field” refers to a location in the event data containing one or more values for a specific data item. As will be described in more detail herein, the fields are defined by extraction rules (e.g., regular expressions) that derive one or more values from the portion of raw machine data in each event that has a particular field specified by an extraction rule. The set of values so produced are semantically-related (such as IP address), even though the raw machine data in each event may be in different formats (e.g., semantically-related values may be in different positions in the events derived from different sources).

[0073] As noted above, the SPLUNK® ENTERPRISE system utilizes a late-binding schema to event data while performing queries on events. One aspect of a late-binding schema is applying “extraction rules” to event data to extract values for specific fields during search time. More specifically, the extraction rules for a field can include one or more instructions that specify how to extract a value for the field from the event data. An extraction rule can generally include any type of instruction for extracting values from data in events. In some cases, an extraction rule comprises a regular expression where a sequence of characters form a search pattern, in which case the rule is referred to as a “regex rule.” The system applies the regex rule to the event data to extract values for associated fields in the event data by searching the event data for the sequence of characters defined in the regex rule.

[0074] In the SPLUNK® ENTERPRISE system, a field extractor may be configured to automatically generate extraction rules for certain field values in the events when the events are being created, indexed, or stored, or possibly at a later time. Alternatively, a user may manually define extraction rules for fields using a variety of techniques. In contrast to a conventional schema for a database system, a late-binding schema is not defined at data ingestion time. Instead, the late-binding schema can be developed on an ongoing basis until the time a query is actually executed. This means that extraction rules for the fields in a query may be provided in the query itself, or may be located during execution of the query. Hence, as a user learns more about the data in the events, the user can continue to refine the late-binding schema by adding new fields, deleting fields, or modifying the field extraction rules for use the next time the schema is used by the system. Because the SPLUNK® ENTERPRISE system maintains the underlying raw data and uses late-binding schema for searching the raw data, it enables a user to continue investigating and learn valuable insights about the raw data.

[0075] In some embodiments, a common field name may be used to reference two or more fields containing equivalent data items, even though the fields may be associated with different types of events that possibly have different data formats and different extraction rules. By enabling a common field name to be used to identify equivalent fields from different types of events generated by disparate data sources, the system facilitates use of a “common information model” (CIM) across the disparate data sources (further discussed with respect to FIG. 5).

2.0. Operating Environment

[0076] FIG. 1 illustrates a networked computer system 100 in which an embodiment may be implemented. Those skilled in the art would understand that FIG. 1 represents one example of a networked computer system and other embodiments may use different arrangements.

[0077] The networked computer system 100 comprises one or more computing devices. These one or more computing devices comprise any combination of hardware and software configured to implement the various logical components described herein. For example, the one or more computing devices may include one or more memories that store instructions for implementing the various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more

memories for storing data structures utilized and manipulated by the various components.

[0078] In an embodiment, one or more client devices 102 are coupled to one or more host devices 106 and a data intake and query system 108 via one or more networks 104.

[0079] Networks 104 broadly represent one or more LANs, WANs, cellular networks (e.g., LTE, HSPA, 3G, and other cellular technologies), and/or networks using any of wired, wireless, terrestrial microwave, or satellite links, and may include the public Internet.

2.1. Host Devices

[0080] In the illustrated embodiment, a system 100 includes one or more host devices 106. Host devices 106 may broadly include any number of computers, virtual machine instances, and/or data centers that are configured to host or execute one or more instances of host applications 114. In general, a host device 106 may be involved, directly or indirectly, in processing requests received from client devices 102. Each host device 106 may comprise, for example, one or more of a network device, a web server, an application server, a database server, etc. A collection of host devices 106 may be configured to implement a network-based service. For example, a provider of a network-based service may configure one or more host devices 106 and host applications 114 (e.g., one or more web servers, application servers, database servers, etc.) to collectively implement the network-based application.

[0081] In general, client devices 102 communicate with one or more host applications 114 to exchange information. The communication between a client device 102 and a host application 114 may, for example, be based on the Hypertext Transfer Protocol (HTTP) or any other network protocol. Content delivered from the host application 114 to a client device 102 may include, for example, HTML documents, media content, etc. The communication between a client device 102 and host application 114 may include sending various requests and receiving data packets. For example, in general, a client device 102 or application running on a client device may initiate communication with a host application 114 by making a request for a specific resource (e.g., based on an HTTP request), and the application server may respond with the requested content stored in one or more response packets.

[0082] In the illustrated embodiment, one or more of host applications 114 may generate various types of performance data during operation, including event logs, network data, sensor data, and other types of machine-generated data. For example, a host application 114 comprising a web server may generate one or more web server logs in which details of interactions between the web server and any number of client devices 102 is recorded. As another example, a host device 106 comprising a router may generate one or more router logs that record information related to network traffic managed by the router. As yet another example, a host application 114 comprising a database server may generate one or more logs that record information related to requests sent from other host applications 114 (e.g., web servers or application servers) for data managed by the database server.

2.2. Client Devices

[0083] Client devices 102 of FIG. 1 represent any computing device capable of interacting with one or more host

devices 106 via a network 104. Examples of client devices 102 may include, without limitation, smart phones, tablet computers, handheld computers, wearable devices, laptop computers, desktop computers, servers, portable media players, gaming devices, and so forth. In general, a client device 102 can provide access to different content, for instance, content provided by one or more host devices 106, etc. Each client device 102 may comprise one or more client applications 110, described in more detail in a separate section hereinafter.

2.3. Client Device Applications

[0084] In an embodiment, each client device 102 may host or execute one or more client applications 110 that are capable of interacting with one or more host devices 106 via one or more networks 104. For instance, a client application 110 may be or comprise a web browser that a user may use to navigate to one or more websites or other resources provided by one or more host devices 106. As another example, a client application 110 may comprise a mobile application or “app.” For example, an operator of a network-based service hosted by one or more host devices 106 may make available one or more mobile apps that enable users of client devices 102 to access various resources of the network-based service. As yet another example, client applications 110 may include background processes that perform various operations without direct interaction from a user. A client application 110 may include a “plug-in” or “extension” to another application, such as a web browser plug-in or extension.

[0085] In an embodiment, a client application 110 may include a monitoring component 112. At a high level, the monitoring component 112 comprises a software component or other logic that facilitates generating performance data related to a client device’s operating state, including monitoring network traffic sent and received from the client device and collecting other device and/or application-specific information. Monitoring component 112 may be an integrated component of a client application 110, a plug-in, an extension, or any other type of add-on component. Monitoring component 112 may also be a stand-alone process.

[0086] In one embodiment, a monitoring component 112 may be created when a client application 110 is developed, for example, by an application developer using a software development kit (SDK). The SDK may include custom monitoring code that can be incorporated into the code implementing a client application 110. When the code is converted to an executable application, the custom code implementing the monitoring functionality can become part of the application itself.

[0087] In some cases, an SDK or other code for implementing the monitoring functionality may be offered by a provider of a data intake and query system, such as a system 108. In such cases, the provider of the system 108 can implement the custom code so that performance data generated by the monitoring functionality is sent to the system 108 to facilitate analysis of the performance data by a developer of the client application or other users.

[0088] In an embodiment, the custom monitoring code may be incorporated into the code of a client application 110 in a number of different ways, such as the insertion of one or more lines in the client application code that call or otherwise invoke the monitoring component 112. As such, a

developer of a client application 110 can add one or more lines of code into the client application 110 to trigger the monitoring component 112 at desired points during execution of the application. Code that triggers the monitoring component may be referred to as a monitor trigger. For instance, a monitor trigger may be included at or near the beginning of the executable code of the client application 110 such that the monitoring component 112 is initiated or triggered as the application is launched, or included at other points in the code that correspond to various actions of the client application, such as sending a network request or displaying a particular interface.

[0089] In an embodiment, the monitoring component 112 may monitor one or more aspects of network traffic sent and/or received by a client application 110. For example, the monitoring component 112 may be configured to monitor data packets transmitted to and/or from one or more host applications 114. Incoming and/or outgoing data packets can be read or examined to identify network data contained within the packets, for example, and other aspects of data packets can be analyzed to determine a number of network performance statistics. Monitoring network traffic may enable information to be gathered particular to the network performance associated with a client application 110 or set of applications.

[0090] In an embodiment, network performance data refers to any type of data that indicates information about the network and/or network performance. Network performance data may include, for instance, a URL requested, a connection type (e.g., HTTP, HTTPS, etc.), a connection start time, a connection end time, an HTTP status code, request length, response length, request headers, response headers, connection status (e.g., completion, response time(s), failure, etc.), and the like. Upon obtaining network performance data indicating performance of the network, the network performance data can be transmitted to a data intake and query system 108 for analysis.

[0091] Upon developing a client application 110 that incorporates a monitoring component 112, the client application 110 can be distributed to client devices 102. Applications generally can be distributed to client devices 102 in any manner, or they can be pre-loaded. In some cases, the application may be distributed to a client device 102 via an application marketplace or other application distribution system. For instance, an application marketplace or other application distribution system might distribute the application to a client device based on a request from the client device to download the application.

[0092] Examples of functionality that enables monitoring performance of a client device are described in U.S. patent application Ser. No. 14/524,748, entitled “UTILIZING PACKET HEADERS TO MONITOR NETWORK TRAFFIC IN ASSOCIATION WITH A CLIENT DEVICE”, filed on 27 Oct. 2014, and which is hereby incorporated by reference in its entirety for all purposes.

[0093] In an embodiment, the monitoring component 112 may also monitor and collect performance data related to one or more aspects of the operational state of a client application 110 and/or client device 102. For example, a monitoring component 112 may be configured to collect device performance information by monitoring one or more client device operations, or by making calls to an operating system and/or one or more other applications executing on a client device 102 for performance information. Device

performance information may include, for instance, a current wireless signal strength of the device, a current connection type and network carrier, current memory performance information, a geographic location of the device, a device orientation, and any other information related to the operational state of the client device.

[0094] In an embodiment, the monitoring component 112 may also monitor and collect other device profile information including, for example, a type of client device, a manufacturer and model of the device, versions of various software applications installed on the device, and so forth.

[0095] In general, a monitoring component 112 may be configured to generate performance data in response to a monitor trigger in the code of a client application 110 or other triggering application event, as described above, and to store the performance data in one or more data records. Each data record, for example, may include a collection of field-value pairs, each field-value pair storing a particular item of performance data in association with a field for the item. For example, a data record generated by a monitoring component 112 may include a “networkLatency” field (not shown in the Figure) in which a value is stored. This field indicates a network latency measurement associated with one or more network requests. The data record may include a “state” field to store a value indicating a state of a network connection, and so forth for any number of aspects of collected performance data.

2.4. Data Server System

[0096] FIG. 2 depicts a block diagram of an exemplary data intake and query system 108, similar to the SPLUNK® ENTERPRISE system. System 108 includes one or more forwarders 204 that receive data from a variety of input data sources 202, and one or more indexers 206 that process and store the data in one or more data stores 208. These forwarders and indexers can comprise separate computer systems, or may alternatively comprise separate processes executing on one or more computer systems.

[0097] Each data source 202 broadly represents a distinct source of data that can be consumed by a system 108. Examples of a data source 202 include, without limitation, data files, directories of files, data sent over a network, event logs, registries, etc.

[0098] During operation, the forwarders 204 identify which indexers 206 receive data collected from a data source 202 and forward the data to the appropriate indexers. Forwarders 204 can also perform operations on the data before forwarding, including removing extraneous data, detecting timestamps in the data, parsing data, indexing data, routing data based on criteria relating to the data being routed, and/or performing other data transformations.

[0099] In an embodiment, a forwarder 204 may comprise a service accessible to client devices 102 and host devices 106 via a network 104. For example, one type of forwarder 204 may be capable of consuming vast amounts of real-time data from a potentially large number of client devices 102 and/or host devices 106. The forwarder 204 may, for example, comprise a computing device which implements multiple data pipelines or “queues” to handle forwarding of network data to indexers 206. A forwarder 204 may also perform many of the functions that are performed by an indexer. For example, a forwarder 204 may perform keyword extractions on raw data or parse raw data to create events. A forwarder 204 may generate time stamps for

events. Additionally or alternatively, a forwarder 204 may perform routing of events to indexers. Data store 208 may contain events derived from machine data from a variety of sources all pertaining to the same component in an IT environment, and this data may be produced by the machine in question or by other components in the IT environment.

2.5. Data Ingestion

[0100] FIG. 3 depicts a flow chart illustrating an example data flow performed by Data Intake and Query system 108, in accordance with the disclosed embodiments. The data flow illustrated in FIG. 3 is provided for illustrative purposes only; those skilled in the art would understand that one or more of the steps of the processes illustrated in FIG. 3 may be removed or the ordering of the steps may be changed. Furthermore, for the purposes of illustrating a clear example, one or more particular system components are described in the context of performing various operations during each of the data flow stages. For example, a forwarder is described as receiving and processing data during an input phase; an indexer is described as parsing and indexing data during parsing and indexing phases; and a search head is described as performing a search query during a search phase. However, other system arrangements and distributions of the processing steps across system components may be used.

2.5.1. Input

[0101] At block 302, a forwarder receives data from an input source, such as a data source 202 shown in FIG. 2. A forwarder initially may receive the data as a raw data stream generated by the input source. For example, a forwarder may receive a data stream from a log file generated by an application server, from a stream of network data from a network device, or from any other source of data. In one embodiment, a forwarder receives the raw data and may segment the data stream into “blocks”, or “buckets,” possibly of a uniform data size, to facilitate subsequent processing steps.

[0102] At block 304, a forwarder or other system component annotates each block generated from the raw data with one or more metadata fields. These metadata fields may, for example, provide information related to the data block as a whole and may apply to each event that is subsequently derived from the data in the data block. For example, the metadata fields may include separate fields specifying each of a host, a source, and a source type related to the data block. A host field may contain a value identifying a host name or IP address of a device that generated the data. A source field may contain a value identifying a source of the data, such as a pathname of a file or a protocol and port related to received network data. A source type field may contain a value specifying a particular source type label for the data. Additional metadata fields may also be included during the input phase, such as a character encoding of the data, if known, and possibly other values that provide information relevant to later processing steps. In an embodiment, a forwarder forwards the annotated data blocks to another system component (typically an indexer) for further processing.

[0103] The SPLUNK® ENTERPRISE system allows forwarding of data from one SPLUNK® ENTERPRISE instance to another, or even to a third-party system.

SPLUNK® ENTERPRISE system can employ different types of forwarders in a configuration.

[0104] In an embodiment, a forwarder may contain the essential components needed to forward data. It can gather data from a variety of inputs and forward the data to a SPLUNK® ENTERPRISE server for indexing and searching. It also can tag metadata (e.g., source, source type, host, etc.).

[0105] Additionally or optionally, in an embodiment, a forwarder has the capabilities of the aforementioned forwarder as well as additional capabilities. The forwarder can parse data before forwarding the data (e.g., associate a time stamp with a portion of data and create an event, etc.) and can route data based on criteria such as source or type of event. It can also index data locally while forwarding the data to another indexer.

2.5.2 Parsing

[0106] At block 306, an indexer receives data blocks from a forwarder and parses the data to organize the data into events. In an embodiment, to organize the data into events, an indexer may determine a source type associated with each data block (e.g., by extracting a source type label from the metadata fields associated with the data block, etc.) and refer to a source type configuration corresponding to the identified source type. The source type definition may include one or more properties that indicate to the indexer to automatically determine the boundaries of events within the data. In general, these properties may include regular expression-based rules or delimiter rules where, for example, event boundaries may be indicated by predefined characters or character strings. These predefined characters may include punctuation marks or other special characters including, for example, carriage returns, tabs, spaces, line breaks, etc. If a source type for the data is unknown to the indexer, an indexer may infer a source type for the data by examining the structure of the data. Then, it can apply an inferred source type definition to the data to create the events.

[0107] At block 308, the indexer determines a timestamp for each event. Similar to the process for creating events, an indexer may again refer to a source type definition associated with the data to locate one or more properties that indicate instructions for determining a timestamp for each event. The properties may, for example, instruct an indexer to extract a time value from a portion of data in the event, to interpolate time values based on timestamps associated with temporally proximate events, to create a timestamp based on a time the event data was received or generated, to use the timestamp of a previous event, or use any other rules for determining timestamps.

[0108] At block 310, the indexer associates with each event one or more metadata fields including a field containing the timestamp (in some embodiments, a timestamp may be included in the metadata fields) determined for the event. These metadata fields may include a number of “default fields” that are associated with all events, and may also include one or more custom fields as defined by a user. Similar to the metadata fields associated with the data blocks at block 304, the default metadata fields associated with each event may include a host, source, and source type field including or in addition to a field storing the timestamp.

[0109] At block 312, an indexer may optionally apply one or more transformations to data included in the events created at block 306. For example, such transformations can

include removing a portion of an event (e.g., a portion used to define event boundaries, extraneous characters from the event, other extraneous text, etc.), masking a portion of an event (e.g., masking a credit card number), removing redundant portions of an event, etc. The transformations applied to event data may, for example, be specified in one or more configuration files and referenced by one or more source type definitions.

2.5.3. Indexing

[0110] At blocks 314 and 316, an indexer can optionally generate a keyword index to facilitate fast keyword searching for event data. To build a keyword index, at block 314, the indexer identifies a set of keywords in each event. At block 316, the indexer includes the identified keywords in an index, which associates each stored keyword with reference pointers to events containing that keyword (or to locations within events where that keyword is located, other location identifiers, etc.). When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword.

[0111] In some embodiments, the keyword index may include entries for name-value pairs found in events, where a name-value pair can include a pair of keywords connected by a symbol, such as an equals sign or colon. This way, events containing these name-value pairs can be quickly located. In some embodiments, fields can automatically be generated for some or all of the name-value pairs at the time of indexing. For example, if the string “dest=10.0.1.2” is found in an event, a field named “dest” may be created for the event, and assigned a value of “10.0.1.2”.

[0112] At block 318, the indexer stores the events with an associated timestamp in a data store 208. Timestamps enable a user to search for events based on a time range. In one embodiment, the stored events are organized into “buckets,” where each bucket stores events associated with a specific time range based on the timestamps associated with each event. This may not only improve time-based searching, but also allows for events with recent timestamps, which may have a higher likelihood of being accessed, to be stored in a faster memory to facilitate faster retrieval. For example, buckets containing the most recent events can be stored in flash memory rather than on a hard disk.

[0113] Each indexer 206 may be responsible for storing and searching a subset of the events contained in a corresponding data store 208. By distributing events among the indexers and data stores, the indexers can analyze events for a query in parallel. For example, using map-reduce techniques, each indexer returns partial responses for a subset of events to a search head that combines the results to produce an answer for the query. By storing events in buckets for specific time ranges, an indexer may further optimize data retrieval process by searching buckets corresponding to time ranges that are relevant to a query.

[0114] Moreover, events and buckets can also be replicated across different indexers and data stores to facilitate high availability and disaster recovery as described in U.S. patent application Ser. No. 14/266,812, entitled “SITE-BASED SEARCH AFFINITY”, filed on 30 Apr. 2014, and in U.S. patent application Ser. No. 14/266,817, entitled “MULTI-SITE CLUSTERING”, also filed on 30 Apr. 2014, each of which is hereby incorporated by reference in its entirety for all purposes.

2.6. Query Processing

[0115] FIG. 4 is a flow diagram that illustrates an exemplary process that a search head and one or more indexers may perform during a search query. At block 402, a search head receives a search query from a client. At block 404, the search head analyzes the search query to determine what portion(s) of the query can be delegated to indexers and what portions of the query can be executed locally by the search head. At block 406, the search head distributes the determined portions of the query to the appropriate indexers. In an embodiment, a search head cluster may take the place of an independent search head where each search head in the search head cluster coordinates with peer search heads in the search head cluster to schedule jobs, replicate search results, update configurations, fulfill search requests, etc. In an embodiment, the search head (or each search head) communicates with a master node (also known as a cluster master, not shown in Fig.) that provides the search head with a list of indexers to which the search head can distribute the determined portions of the query. The master node maintains a list of active indexers and can also designate which indexers may have responsibility for responding to queries over certain sets of events. A search head may communicate with the master node before the search head distributes queries to indexers to discover the addresses of active indexers.

[0116] At block 408, the indexers to which the query was distributed, search data stores associated with them for events that are responsive to the query. To determine which events are responsive to the query, the indexer searches for events that match the criteria specified in the query. These criteria can include matching keywords or specific values for certain fields. The searching operations at block 408 may use the late-binding schema to extract values for specified fields from events at the time the query is processed. In an embodiment, one or more rules for extracting field values may be specified as part of a source type definition. The indexers may then either send the relevant events back to the search head, or use the events to determine a partial result, and send the partial result back to the search head.

[0117] At block 410, the search head combines the partial results and/or events received from the indexers to produce a final result for the query. This final result may comprise different types of data depending on what the query requested. For example, the results can include a listing of matching events returned by the query, or some type of visualization of the data from the returned events. In another example, the final result can include one or more calculated values derived from the matching events.

[0118] The results generated by the system 108 can be returned to a client using different techniques. For example, one technique streams results or relevant events back to a client in real-time as they are identified. Another technique waits to report the results to the client until a complete set of results (which may include a set of relevant events or a result based on relevant events) is ready to return to the client. Yet another technique streams interim results or relevant events back to the client in real-time until a complete set of results is ready, and then returns the complete set of results to the client. In another technique, certain results are stored as “search jobs” and the client may retrieve the results by referring the search jobs.

[0119] The search head can also perform various operations to make the search more efficient. For example, before

the search head begins execution of a query, the search head can determine a time range for the query and a set of common keywords that all matching events include. The search head may then use these parameters to query the indexers to obtain a superset of the eventual results. Then, during a filtering stage, the search head can perform field-extraction operations on the superset to produce a reduced set of search results. This speeds up queries that are performed on a periodic basis.

2.7. Field Extraction

[0120] The search head **210** allows users to search and visualize event data extracted from raw machine data received from homogenous data sources. It also allows users to search and visualize event data extracted from raw machine data received from heterogeneous data sources. The search head **210** includes various mechanisms, which may additionally reside in an indexer **206**, for processing a query. Splunk Processing Language (SPL), used in conjunction with the SPLUNK® ENTERPRISE system, can be utilized to make a query. SPL is a pipelined search language in which a set of inputs is operated on by a first command in a command line, and then a subsequent command following the pipe symbol “|” operates on the results produced by the first command, and so on for additional commands. Other query languages, such as the Structured Query Language (“SQL”), can be used to create a query.

[0121] In response to receiving the search query, search head **210** uses extraction rules to extract values for the fields associated with a field or fields in the event data being searched. The search head **210** obtains extraction rules that specify how to extract a value for certain fields from an event. Extraction rules can comprise regex rules that specify how to extract values for the relevant fields. In addition to specifying how to extract field values, the extraction rules may also include instructions for deriving a field value by performing a function on a character string or value retrieved by the extraction rule. For example, a transformation rule may truncate a character string, or convert the character string into a different data format. In some cases, the query itself can specify one or more extraction rules.

[0122] The search head **210** can apply the extraction rules to event data that it receives from indexers **206**. Indexers **206** may apply the extraction rules to events in an associated data store **208**. Extraction rules can be applied to all the events in a data store, or to a subset of the events that have been filtered based on some criteria (e.g., event time stamp values, etc.). Extraction rules can be used to extract one or more values for a field from events by parsing the event data and examining the event data for one or more patterns of characters, numbers, delimiters, etc., that indicate where the field begins and, optionally, ends.

[0123] FIG. 5 illustrates an example of raw machine data received from disparate data sources. In this example, a user submits an order for merchandise using a vendor's shopping application program **501** running on the user's system. In this example, the order was not delivered to the vendor's server due to a resource exception at the destination server that is detected by the middleware code **502**. The user then sends a message to the customer support **503** to complain about the order failing to complete. The three systems **501**, **502**, and **503** are disparate systems that do not have a common logging format. The order application **501** sends log data **504** to the SPLUNK® ENTERPRISE system in one

format, the middleware code **502** sends error log data **505** in a second format, and the support server **503** sends log data **506** in a third format.

[0124] Using the log data received at one or more indexers **206** from the three systems the vendor can uniquely obtain an insight into user activity, user experience, and system behavior. The search head **210** allows the vendor's administrator to search the log data from the three systems that one or more indexers **206** are responsible for searching, thereby obtaining correlated information, such as the order number and corresponding customer ID number of the person placing the order. The system also allows the administrator to see a visualization of related events via a user interface. The administrator can query the search head **210** for customer ID field value matches across the log data from the three systems that are stored at the one or more indexers **206**. The customer ID field value exists in the data gathered from the three systems, but the customer ID field value may be located in different areas of the data given differences in the architecture of the systems—there is a semantic relationship between the customer ID field values generated by the three systems. The search head **210** requests event data from the one or more indexers **206** to gather relevant event data from the three systems. It then applies extraction rules to the event data in order to extract field values that it can correlate. The search head may apply a different extraction rule to each set of events from each system when the event data format differs among systems. In this example, the user interface can display to the administrator the event data corresponding to the common customer ID field values **507**, **508**, and **509**, thereby providing the administrator with insight into a customer's experience.

[0125] Note that query results can be returned to a client, a search head, or any other system component for further processing. In general, query results may include a set of one or more events, a set of one or more values obtained from the events, a subset of the values, statistics calculated based on the values, a report containing the values, or a visualization, such as a graph or chart, generated from the values.

2.8. Example Search Screen

[0126] FIG. 6A illustrates an example search screen **600** in accordance with the disclosed embodiments. Search screen **600** includes a search bar **602** that accepts user input in the form of a search string. It also includes a time range picker **612** that enables the user to specify a time range for the search. For “historical searches” the user can select a specific time range, or alternatively a relative time range, such as “today,” “yesterday” or “last week.” For “real-time searches,” the user can select the size of a preceding time window to search for real-time events. Search screen **600** also initially displays a “data summary” dialog as is illustrated in FIG. 6B that enables the user to select different sources for the event data, such as by selecting specific hosts and log files.

[0127] After the search is executed, the search screen **600** in FIG. 6A can display the results through search results tabs **604**, wherein search results tabs **604** includes: an “events tab” that displays various information about events returned by the search; a “statistics tab” that displays statistics about the search results; and a “visualization tab” that displays various visualizations of the search results. The events tab illustrated in FIG. 6A displays a timeline graph **605** that graphically illustrates the number of events that occurred in

one-hour intervals over the selected time range. It also displays an events list 608 that enables a user to view the raw data in each of the returned events. It additionally displays a fields sidebar 606 that includes statistics about occurrences of specific fields in the returned events, including “selected fields” that are pre-selected by the user, and “interesting fields” that are automatically selected by the system based on pre-specified criteria.

2.9. Data Models

[0128] A data model is a hierarchically structured search-time mapping of semantic knowledge about one or more datasets. It encodes the domain knowledge necessary to build a variety of specialized searches of those datasets. Those searches, in turn, can be used to generate reports.

[0129] A data model is composed of one or more “objects” (or “data model objects”) that define or otherwise correspond to a specific set of data.

[0130] Objects in data models can be arranged hierarchically in parent/child relationships. Each child object represents a subset of the dataset covered by its parent object. The top-level objects in data models are collectively referred to as “root objects.”

[0131] Child objects have inheritance. Data model objects are defined by characteristics that mostly break down into constraints and attributes. Child objects inherit constraints and attributes from their parent objects and have additional constraints and attributes of their own. Child objects provide a way of filtering events from parent objects. Because a child object always provides an additional constraint in addition to the constraints it has inherited from its parent object, the dataset it represents is always a subset of the dataset that its parent represents.

[0132] For example, a first data model object may define a broad set of data pertaining to e-mail activity generally, and another data model object may define specific datasets within the broad dataset, such as a subset of the e-mail data pertaining specifically to e-mails sent. Examples of data models can include electronic mail, authentication, databases, intrusion detection, malware, application state, alerts, compute inventory, network sessions, network traffic, performance, audits, updates, vulnerabilities, etc. Data models and their objects can be designed by knowledge managers in an organization, and they can enable downstream users to quickly focus on a specific set of data. For example, a user can simply select an “e-mail activity” data model object to access a dataset relating to e-mails generally (e.g., sent or received), or select an “e-mails sent” data model object (or data sub-model object) to access a dataset relating to e-mails sent.

[0133] A data model object may be defined by (1) a set of search constraints, and (2) a set of fields. Thus, a data model object can be used to quickly search data to identify a set of events and to identify a set of fields to be associated with the set of events. For example, an “e-mails sent” data model object may specify a search for events relating to e-mails that have been sent, and specify a set of fields that are associated with the events. Thus, a user can retrieve and use the “e-mails sent” data model object to quickly search source data for events relating to sent e-mails, and may be provided with a listing of the set of fields relevant to the events in a user interface screen.

[0134] A child of the parent data model may be defined by a search (typically a narrower search) that produces a subset

of the events that would be produced by the parent data model’s search. The child’s set of fields can include a subset of the set of fields of the parent data model and/or additional fields. Data model objects that reference the subsets can be arranged in a hierarchical manner, so that child subsets of events are proper subsets of their parents. A user iteratively applies a model development tool (not shown in Fig.) to prepare a query that defines a subset of events and assigns an object name to that subset. A child subset is created by further limiting a query that generated a parent subset. A late-binding schema of field extraction rules is associated with each object or subset in the data model.

[0135] Data definitions in associated schemas can be taken from the common information model (CIM) or can be devised for a particular schema and optionally added to the CIM. Child objects inherit fields from parents and can include fields not present in parents. A model developer can select fewer extraction rules than are available for the sources returned by the query that defines events belonging to a model. Selecting a limited set of extraction rules can be a tool for simplifying and focusing the data model, while allowing a user flexibility to explore the data subset. Development of a data model is further explained in U.S. Pat. Nos. 8,788,525 and 8,788,526, both entitled “DATA MODEL FOR MACHINE DATA FOR SEMANTIC SEARCH”, both issued on 22 Jul. 2014, U.S. Pat. No. 8,983,994, entitled “GENERATION OF A DATA MODEL FOR SEARCHING MACHINE DATA”, issued on 17 Mar. 2015, U.S. patent application Ser. No. 14/611,232, entitled “GENERATION OF A DATA MODEL APPLIED TO QUERIES”, filed on 31 Jan. 2015, and U.S. patent application Ser. No. 14/815,884, entitled “GENERATION OF A DATA MODEL APPLIED TO OBJECT QUERIES”, filed on 31 Jul. 2015, each of which is hereby incorporated by reference in its entirety for all purposes. See, also, Knowledge Manager Manual, Build a Data Model, Splunk Enterprise 6.1.3 pp. 150-204 (Aug. 25, 2014).

[0136] A data model can also include reports. One or more report formats can be associated with a particular data model and be made available to run against the data model. A user can use child objects to design reports with object datasets that already have extraneous data pre-filtered out. In an embodiment, the data intake and query system 108 provides the user with the ability to produce reports (e.g., a table, chart, visualization, etc.) without having to enter SPL, SQL, or other query language terms into a search screen. Data models are used as the basis for the search feature.

[0137] Data models may be selected in a report generation interface. The report generator supports drag-and-drop organization of fields to be summarized in a report. When a model is selected, the fields with available extraction rules are made available for use in the report. The user may refine and/or filter search results to produce more precise reports. The user may select some fields for organizing the report and select other fields for providing detail according to the report organization. For example, “region” and “salesperson” are fields used for organizing the report and sales data can be summarized (subtotaled and totaled) within this organization. The report generator allows the user to specify one or more fields within events and apply statistical analysis on values extracted from the specified one or more fields. The report generator may aggregate search results across sets of events and generate statistics based on aggregated search results. Building reports using the report generation inter-

face is further explained in U.S. patent application Ser. No. 14/503,335, entitled “GENERATING REPORTS FROM UNSTRUCTURED DATA”, filed on 30 Sep. 2014, and published as US Patent Application Publication 2015-0019537A1 on Jan. 15, 2015, and which is hereby incorporated by reference in its entirety for all purposes, and in Pivot Manual, Splunk Enterprise 6.1.3 (Aug. 4, 2014). Data visualizations also can be generated in a variety of formats, by reference to the data model. Reports, data visualizations, and data model objects can be saved and associated with the data model for future use. The data model object may be used to perform searches of other data.

2.10. Acceleration Technique

[0138] The above-described system provides significant flexibility by enabling a user to analyze massive quantities of minimally processed data “on the fly” at search time instead of storing pre-specified portions of the data in a database at ingestion time. This flexibility enables a user to see valuable insights, correlate data, and perform subsequent queries to examine interesting aspects of the data that may not have been apparent at ingestion time.

[0139] However, performing extraction and analysis operations at search time can involve a large amount of data and require a large number of computational operations, which can cause delays in processing the queries. Advantageously, SPLUNK® ENTERPRISE system employs a number of unique acceleration techniques that have been developed to speed up analysis operations performed at search time. These techniques include: (1) performing search operations in parallel across multiple indexers; (2) using a keyword index; (3) using a high performance analytics store; and (4) accelerating the process of generating reports. These novel techniques are described in more detail below.

2.10.1 Aggregation Technique

[0140] To facilitate faster query processing, a query can be structured such that multiple indexers perform the query in parallel, while aggregation of search results from the multiple indexers is performed locally at the search head. For example, FIG. 7 illustrates how a search query 802 received from a client at a search head 210 can split into two phases, including: (1) subtasks 804 (e.g., data retrieval or simple filtering) that may be performed in parallel by indexers 206 for execution, and (2) a search results aggregation operation 806 to be executed by the search head when the results are ultimately collected from the indexers.

[0141] During operation, upon receiving search query 802, a search head 210 determines that a portion of the operations involved with the search query may be performed locally by the search head. The search head modifies search query 802 by substituting “stats” (create aggregate statistics over results sets received from the indexers at the search head) with “prestats” (create statistics by the indexer from local results set) to produce search query 804, and then distributes search query 804 to distributed indexers, which are also referred to as “search peers.” Note that search queries may generally specify search criteria or operations to be performed on events that meet the search criteria. Search queries may also specify field names, as well as search criteria for the values in the fields or operations to be performed on the values in the fields. Moreover, the search head may distribute the full search query to the search peers

as illustrated in FIG. 4, or may alternatively distribute a modified version (e.g., a more restricted version) of the search query to the search peers. In this example, the indexers are responsible for producing the results and sending them to the search head. After the indexers return the results to the search head, the search head aggregates the received results 806 to form a single search result set. By executing the query in this manner, the system effectively distributes the computational operations across the indexers while minimizing data transfers.

2.10.2 Keyword Index

[0142] As described above with reference to the flow charts in FIG. 3 and FIG. 4, data intake and query system 108 can construct and maintain one or more keyword indices to quickly identify events containing specific keywords. This technique can greatly speed up the processing of queries involving specific keywords. As mentioned above, to build a keyword index, an indexer first identifies a set of keywords. Then, the indexer includes the identified keywords in an index, which associates each stored keyword with references to events containing that keyword, or to locations within events where that keyword is located. When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword.

2.10.3. High Performance Analytics Store

[0143] To speed up certain types of queries, some embodiments of system 108 create a high performance analytics store, which is referred to as a “summarization table,” that contains entries for specific field-value pairs. Each of these entries keeps track of instances of a specific value in a specific field in the event data and includes references to events containing the specific value in the specific field. For example, an example entry in a summarization table can keep track of occurrences of the value “94107” in a “ZIP code” field of a set of events and the entry includes references to all of the events that contain the value “94107” in the ZIP code field. This optimization technique enables the system to quickly process queries that seek to determine how many events have a particular value for a particular field. To this end, the system can examine the entry in the summarization table to count instances of the specific value in the field without having to go through the individual events or perform data extractions at search time. Also, if the system needs to process all events that have a specific field-value combination, the system can use the references in the summarization table entry to directly access the events to extract further information without having to search all of the events to find the specific field-value combination at search time.

[0144] In some embodiments, the system maintains a separate summarization table for each of the above-described time-specific buckets that stores events for a specific time range. A bucket-specific summarization table includes entries for specific field-value combinations that occur in events in the specific bucket. Alternatively, the system can maintain a separate summarization table for each indexer. The indexer-specific summarization table includes entries for the events in a data store that are managed by the specific indexer. Indexer-specific summarization tables may also be bucket-specific.

[0145] The summarization table can be populated by running a periodic query that scans a set of events to find instances of a specific field-value combination, or alternatively instances of all field-value combinations for a specific field. A periodic query can be initiated by a user, or can be scheduled to occur automatically at specific time intervals. A periodic query can also be automatically launched in response to a query that asks for a specific field-value combination.

[0146] In some cases, when the summarization tables may not cover all of the events that are relevant to a query, the system can use the summarization tables to obtain partial results for the events that are covered by summarization tables, but may also have to search through other events that are not covered by the summarization tables to produce additional results. These additional results can then be combined with the partial results to produce a final set of results for the query. The summarization table and associated techniques are described in more detail in U.S. Pat. No. 8,682,925, entitled "Distributed High Performance Analytics Store", issued on 25 Mar. 2014, U.S. patent application Ser. No. 14/170,159, entitled "SUPPLEMENTING A HIGH PERFORMANCE ANALYTICS STORE WITH EVALUATION OF INDIVIDUAL EVENTS TO RESPOND TO AN EVENT QUERY", filed on 31 Jan. 2014 and issued as U.S. Pat. No. 9,128,985 on Sep. 8, 2015, and U.S. patent application Ser. No. 14/815,973, entitled "STORAGE MEDIUM AND CONTROL DEVICE", filed on 21 Feb. 2014, each of which is hereby incorporated by reference in its entirety.

2.10.4. Accelerating Report Generation

[0147] In some embodiments, a data server system such as the SPLUNK® ENTERPRISE system can accelerate the process of periodically generating updated reports based on query results. To accelerate this process, a summarization engine automatically examines the query to determine whether generation of updated reports can be accelerated by creating intermediate summaries. If reports can be accelerated, the summarization engine periodically generates a summary covering data obtained during a latest non-overlapping time period. For example, where the query seeks events meeting a specified criteria, a summary for the time period includes only events within the time period that meet the specified criteria. Similarly, if the query seeks statistics calculated from the events, such as the number of events that match the specified criteria, then the summary for the time period includes the number of events in the period that match the specified criteria.

[0148] In addition to the creation of the summaries, the summarization engine schedules the periodic updating of the report associated with the query. During each scheduled report update, the query engine determines whether intermediate summaries have been generated covering portions of the time period covered by the report update. If so, then the report is generated based on the information contained in the summaries. Also, if additional event data has been received and has not yet been summarized, and is required to generate the complete report, the query can be run on this additional event data. Then, the results returned by this query on the additional event data, along with the partial results obtained from the intermediate summaries, can be combined to generate the updated report. This process is repeated each time the report is updated. Alternatively, if the system stores events in buckets covering specific time

ranges, then the summaries can be generated on a bucket-by-bucket basis. Note that producing intermediate summaries can save the work involved in re-running the query for previous time periods, so advantageously only the newer event data needs to be processed while generating an updated report. These report acceleration techniques are described in more detail in U.S. Pat. No. 8,589,403, entitled "Compressed Journaling In Event Tracking Files For Metadata Recovery And Replication", issued on 19 Nov. 2013, U.S. Pat. No. 8,412,696, entitled "Real Time Searching And Reporting", issued on 2 Apr. 2011, and U.S. Pat. Nos. 8,589,375 and 8,589,432, both also entitled "REAL TIME SEARCHING AND REPORTING", both issued on 19 Nov. 2013, each of which is hereby incorporated by reference in its entirety.

2.13. Cloud-Based System Overview

[0149] The example data intake and query system 108 described in reference to FIG. 2 comprises several system components, including one or more forwarders, indexers, and search heads. In some environments, a user of a data intake and query system 108 may install and configure, on computing devices owned and operated by the user, one or more software applications that implement some or all of these system components. For example, a user may install a software application on server computers owned by the user and configure each server to operate as one or more of a forwarder, an indexer, a search head, etc. This arrangement generally may be referred to as an "on-premises" solution. That is, the system 108 is installed and operates on computing devices directly controlled by the user of the system. Some users may prefer an on-premises solution because it may provide a greater level of control over the configuration of certain aspects of the system (e.g., security, privacy, standards, controls, etc.). However, other users may instead prefer an arrangement in which the user is not directly responsible for providing and managing the computing devices upon which various components of system 108 operate.

[0150] In one embodiment, to provide an alternative to an entirely on-premises environment for system 108, one or more of the components of a data intake and query system instead may be provided as a cloud-based service. In this context, a cloud-based service refers to a service hosted by one or more computing resources that are accessible to end users over a network, for example, by using a web browser or other application on a client device to interface with the remote computing resources. For example, a service provider may provide a cloud-based data intake and query system by managing computing resources configured to implement various aspects of the system (e.g., forwarders, indexers, search heads, etc.) and by providing access to the system to end users via a network. Typically, a user may pay a subscription or other fee to use such a service. Each subscribing user of the cloud-based service may be provided with an account that enables the user to configure a customized cloud-based system based on the user's preferences.

[0151] FIG. 8 illustrates a block diagram of an example cloud-based data intake and query system. Similar to the system of FIG. 2, the networked computer system 1000 includes input data sources 202 and forwarders 204. These input data sources and forwarders may be in a subscriber's private computing environment. Alternatively, they might be directly managed by the service provider as part of the cloud

service. In the example system **1000**, one or more forwarders **204** and client devices **1002** are coupled to a cloud-based data intake and query system **1006** via one or more networks **1004**. Network **1004** broadly represents one or more LANs, WANs, cellular networks, intranetworks, internetworks, etc., using any of wired, wireless, terrestrial microwave, satellite links, etc., and may include the public Internet, and is used by client devices **1002** and forwarders **204** to access the system **1006**. Similar to the system of **108**, each of the forwarders **204** may be configured to receive data from an input source and to forward the data to other components of the system **1006** for further processing.

[0152] In an embodiment, a cloud-based data intake and query system **1006** may comprise a plurality of system instances **1008**. In general, each system instance **1008** may include one or more computing resources managed by a provider of the cloud-based system **1006** made available to a particular subscriber. The computing resources comprising a system instance **1008** may, for example, include one or more servers or other devices configured to implement one or more forwarders, indexers, search heads, and other components of a data intake and query system, similar to system **108**. As indicated above, a subscriber may use a web browser or other application of a client device **1002** to access a web portal or other interface that enables the subscriber to configure an instance **1008**.

[0153] Providing a data intake and query system as described in reference to system **108** as a cloud-based service presents a number of challenges. Each of the components of a system **108** (e.g., forwarders, indexers and search heads) may at times refer to various configuration files stored locally at each component. These configuration files typically may involve some level of user configuration to accommodate particular types of data a user desires to analyze and to account for other user preferences. However, in a cloud-based service context, users typically may not have direct access to the underlying computing resources implementing the various system components (e.g., the computing resources comprising each system instance **1008**) and may desire to make such configurations indirectly, for example, using one or more web-based interfaces. Thus, the techniques and systems described herein for providing user interfaces that enable a user to configure source type definitions are applicable to both on-premises and cloud-based service contexts, or some combination thereof (e.g., a hybrid system where both an on-premises environment such as SPLUNK® ENTERPRISE and a cloud-based environment such as SPLUNK CLOUD™ are centrally visible).

2.14. Searching Externally Archived Data

[0154] FIG. 9 shows a block diagram of an example of a data intake and query system **108** that provides transparent search facilities for data systems that are external to the data intake and query system. Such facilities are available in the HUNK® system provided by Splunk Inc. of San Francisco, California. HUNK® represents an analytics platform that enables business and IT teams to rapidly explore, analyze, and visualize data in Hadoop and NoSQL data stores.

[0155] The search head **210** of the data intake and query system receives search requests from one or more client devices **1104** over network connections **1120**. As discussed above, the data intake and query system **108** may reside in an enterprise location, in the cloud, etc. FIG. 9 illustrates that multiple client devices **1104a**, **1104b**, . . . , **1104n** may

communicate with the data intake and query system **108**. The client devices **1104** may communicate with the data intake and query system using a variety of connections. For example, one client device in FIG. 9 is illustrated as communicating over an Internet (Web) protocol, another client device is illustrated as communicating via a command line interface, and another client device is illustrated as communicating via a system developer kit (SDK).

[0156] The search head **210** analyzes the received search request to identify request parameters. If a search request received from one of the client devices **1104** references an index maintained by the data intake and query system, then the search head **210** connects to one or more indexers **206** of the data intake and query system for the index referenced in the request parameters. That is, if the request parameters of the search request reference an index, then the search head accesses the data in the index via the indexer. The data intake and query system **108** may include one or more indexers **206**, depending on system access resources and requirements. As described further below, the indexers **206** retrieve data from their respective local data stores **208** as specified in the search request. The indexers and their respective data stores can comprise one or more storage devices and typically reside on the same system, though they may be connected via a local network connection.

[0157] If the request parameters of the received search request reference an external data collection, which is not accessible to the indexers **206** or under the management of the data intake and query system, then the search head **210** can access the external data collection through an External Result Provider (ERP) process **1110**. An external data collection may be referred to as a “virtual index” (plural, “virtual indices”). An ERP process provides an interface through which the search head **210** may access virtual indices.

[0158] Thus, a search reference to an index of the system relates to a locally stored and managed data collection. In contrast, a search reference to a virtual index relates to an externally stored and managed data collection, which the search head may access through one or more ERP processes **1110**, **1112**. FIG. 9 shows two ERP processes **1110**, **1112** that connect to respective remote (external) virtual indices, which are indicated as a Hadoop or another system **1114** (e.g., Amazon S3, Amazon EMR, other Hadoop Compatible File Systems (HCFS), etc.) and a relational database management system (RDB MS) **1116**. Other virtual indices may include other file organizations and protocols, such as Structured Query Language (SQL) and the like. The ellipses between the ERP processes **1110**, **1112** indicate optional additional ERP processes of the data intake and query system **108**. An ERP process may be a computer process that is initiated or spawned by the search head **210** and is executed by the search data intake and query system **108**. Alternatively or additionally, an ERP process may be a process spawned by the search head **210** on the same or different host system as the search head **210** resides.

[0159] The search head **210** may spawn a single ERP process in response to multiple virtual indices referenced in a search request, or the search head may spawn different ERP processes for different virtual indices. Generally, virtual indices that share common data configurations or protocols may share ERP processes. For example, all search query references to a Hadoop file system may be processed by the same ERP process, if the ERP process is suitably configured.

Likewise, all search query references to an SQL database may be processed by the same ERP process. In addition, the search head may provide a common ERP process for common external data source types (e.g., a common vendor may utilize a common ERP process, even if the vendor includes different data storage system types, such as Hadoop and SQL). Common indexing schemes also may be handled by common ERP processes, such as flat text files or Weblog files.

[0160] The search head **210** determines the number of ERP processes to be initiated via the use of configuration parameters that are included in a search request message. Generally, there is a one-to-many relationship between an external results provider “family” and ERP processes. There is also a one-to-many relationship between an ERP process and corresponding virtual indices that are referred to in a search request. For example, using RDBMS, assume two independent instances of such a system by one vendor, such as one RDBMS for production and another RDBMS used for development. In such a situation, it is likely preferable (but optional) to use two ERP processes to maintain the independent operation as between production and development data. Both of the ERPs, however, will belong to the same family, because the two RDBMS system types are from the same vendor.

[0161] The ERP processes **1110**, **1112** receive a search request from the search head **210**. The search head may optimize the received search request for execution at the respective external virtual index. Alternatively, the ERP process may receive a search request as a result of analysis performed by the search head or by a different system process. The ERP processes **1110**, **1112** can communicate with the search head **210** via conventional input/output routines (e.g., standard in/standard out, etc.). In this way, the ERP process receives the search request from a client device such that the search request may be efficiently executed at the corresponding external virtual index.

[0162] The ERP processes **1110**, **1112** may be implemented as a process of the data intake and query system. Each ERP process may be provided by the data intake and query system, or may be provided by process or application providers who are independent of the data intake and query system. Each respective ERP process may include an interface application installed at a computer of the external result provider that ensures proper communication between the search support system and the external result provider. The ERP processes **1110**, **1112** generate appropriate search requests in the protocol and syntax of the respective virtual indices **1114**, **1116**, each of which corresponds to the search request received by the search head **210**. Upon receiving search results from their corresponding virtual indices, the respective ERP process passes the result to the search head **210**, which may return or display the results or a processed set of results based on the returned results to the respective client device.

[0163] Client devices **1104** may communicate with the data intake and query system **108** through a network interface **1120**, e.g., one or more LANs, WANs, cellular networks, intranetworks, and/or internetworks using any of wired, wireless, terrestrial microwave, satellite links, etc., and may include the public Internet.

[0164] The analytics platform utilizing the External Result Provider process described in more detail in U.S. Pat. No. 8,738,629, entitled “External Result Provided Process For

RetriEving Data Stored Using A Different Configuration Or Protocol”, issued on 27 May 2014, U.S. Pat. No. 8,738,587, entitled “PROCESSING A SYSTEM SEARCH REQUEST BY RETRIEVING RESULTS FROM BOTH A NATIVE INDEX AND A VIRTUAL INDEX”, issued on 25 Jul. 2013, U.S. patent application Ser. No. 14/266,832, entitled “PROCESSING A SYSTEM SEARCH REQUEST ACROSS DISPARATE DATA COLLECTION SYSTEMS”, filed on 1 May 2014, and U.S. patent application Ser. No. 14/449,144, entitled “PROCESSING A SYSTEM SEARCH REQUEST INCLUDING EXTERNAL DATA SOURCES”, filed on 31 Jul. 2014, each of which is hereby incorporated by reference in its entirety for all purposes.

2.14.1. ERP Process Features

[0165] The ERP processes described above may include two operation modes: a streaming mode and a reporting mode. The ERP processes can operate in streaming mode only, in reporting mode only, or in both modes simultaneously. Operating in both modes simultaneously is referred to as mixed mode operation. In a mixed mode operation, the ERP at some point can stop providing the search head with streaming results and only provide reporting results thereafter, or the search head at some point may start ignoring streaming results it has been using and only use reporting results thereafter.

[0166] The streaming mode returns search results in real time, with minimal processing, in response to the search request. The reporting mode provides results of a search request with processing of the search results prior to providing them to the requesting search head, which in turn provides results to the requesting client device. ERP operation with such multiple modes provides greater performance flexibility with regard to report time, search latency, and resource utilization.

[0167] In a mixed mode operation, both streaming mode and reporting mode are operating simultaneously. The streaming mode results (e.g., the raw data obtained from the external data source) are provided to the search head, which can then process the results data (e.g., break the raw data into events, timestamp it, filter it, etc.) and integrate the results data with the results data from other external data sources, and/or from data stores of the search head. The search head performs such processing and can immediately start returning interim (streaming mode) results to the user at the requesting client device; simultaneously, the search head is waiting for the ERP process to process the data it is retrieving from the external data source as a result of the concurrently executing reporting mode.

[0168] In some instances, the ERP process initially operates in a mixed mode, such that the streaming mode operates to enable the ERP quickly to return interim results (e.g., some of the raw or unprocessed data necessary to respond to a search request) to the search head, enabling the search head to process the interim results and begin providing to the client or search requester interim results that are responsive to the query. Meanwhile, in this mixed mode, the ERP also operates concurrently in reporting mode, processing portions of raw data in a manner responsive to the search query. Upon determining that it has results from the reporting mode available to return to the search head, the ERP may halt processing in the mixed mode at that time (or some later time) by stopping the return of data in streaming mode to the search head and switching to reporting mode only. The ERP

at this point starts sending interim results in reporting mode to the search head, which in turn may then present this processed data responsive to the search request to the client or search requester. Typically the search head switches from using results from the ERP's streaming mode of operation to results from the ERP's reporting mode of operation when the higher bandwidth results from the reporting mode outstrip the amount of data processed by the search head in the streaming mode of ERP operation.

[0169] A reporting mode may have a higher bandwidth because the ERP does not have to spend time transferring data to the search head for processing all the raw data. In addition, the ERP may optionally direct another processor to do the processing.

[0170] The streaming mode of operation does not need to be stopped to gain the higher bandwidth benefits of a reporting mode; the search head could simply stop using the streaming mode results—and start using the reporting mode results—when the bandwidth of the reporting mode has caught up with or exceeded the amount of bandwidth provided by the streaming mode. Thus, a variety of triggers and ways to accomplish a search head's switch from using streaming mode results to using reporting mode results may be appreciated by one skilled in the art.

[0171] The reporting mode can involve the ERP process (or an external system) performing event breaking, time stamping, filtering of events to match the search query request, and calculating statistics on the results. The user can request particular types of data, such as if the search query itself involves types of events, or the search request may ask for statistics on data, such as on events that meet the search request. In either case, the search head understands the query language used in the received query request, which may be a proprietary language. One exemplary query language is Splunk Processing Language (SPL) developed by the assignee of the application, Splunk Inc. The search head typically understands how to use that language to obtain data from the indexers, which store data in a format used by the SPLUNK® Enterprise system.

[0172] The ERP processes support the search head, as the search head is not ordinarily configured to understand the format in which data is stored in external data sources such as Hadoop or SQL data systems. Rather, the ERP process performs that translation from the query submitted in the search support system's native format (e.g., SPL if SPLUNK® ENTERPRISE is used as the search support system) to a search query request format that will be accepted by the corresponding external data system. The external data system typically stores data in a different format from that of the search support system's native index format, and it utilizes a different query language (e.g., SQL or MapReduce, rather than SPL or the like).

[0173] As noted, the ERP process can operate in the streaming mode alone. After the ERP process has performed the translation of the query request and received raw results from the streaming mode, the search head can integrate the returned data with any data obtained from local data sources (e.g., native to the search support system), other external data sources, and other ERP processes (if such operations were required to satisfy the terms of the search query). An advantage of mixed mode operation is that, in addition to streaming mode, the ERP process is also executing concurrently in reporting mode. Thus, the ERP process (rather than the search head) is processing query results (e.g., performing

event breaking, timestamping, filtering, possibly calculating statistics if required to be responsive to the search query request, etc.). It should be apparent to those skilled in the art that additional time is needed for the ERP process to perform the processing in such a configuration. Therefore, the streaming mode will allow the search head to start returning interim results to the user at the client device before the ERP process can complete sufficient processing to start returning any search results. The switchover between streaming and reporting mode happens when the ERP process determines that the switchover is appropriate, such as when the ERP process determines it can begin returning meaningful results from its reporting mode.

[0174] The operation described above illustrates the source of operational latency: streaming mode has low latency (immediate results) and usually has relatively low bandwidth (fewer results can be returned per unit of time). In contrast, the concurrently running reporting mode has relatively high latency (it has to perform a lot more processing before returning any results) and usually has relatively high bandwidth (more results can be processed per unit of time). For example, when the ERP process does begin returning report results, it returns more processed results than in the streaming mode, because, e.g., statistics only need to be calculated to be responsive to the search request. That is, the ERP process doesn't have to take time to first return raw data to the search head. As noted, the ERP process could be configured to operate in streaming mode alone and return just the raw data for the search head to process in a way that is responsive to the search request. Alternatively, the ERP process can be configured to operate in the reporting mode only. Also, the ERP process can be configured to operate in streaming mode and reporting mode concurrently, as described, with the ERP process stopping the transmission of streaming results to the search head when the concurrently running reporting mode has caught up and started providing results. The reporting mode does not require the processing of all raw data that is responsive to the search query request before the ERP process starts returning results; rather, the reporting mode usually performs processing of chunks of events and returns the processing results to the search head for each chunk.

[0175] For example, an ERP process can be configured to merely return the contents of a search result file verbatim, with little or no processing of results. That way, the search head performs all processing (such as parsing byte streams into events, filtering, etc.). The ERP process can be configured to perform additional intelligence, such as analyzing the search request and handling all the computation that a native search indexer process would otherwise perform. In this way, the configured ERP process provides greater flexibility in features while operating according to desired preferences, such as response latency and resource requirements.

2.15 Cloud-Based Architecture

[0176] As shown in the previous figures, various embodiments may refer to a data intake and query system 108 that includes one or more of a search head 210, an indexer 206, and a forwarder 204. In other implementations, data intake and query system 108 may have a different architecture, but may carry out indexing and searching in a way that is indistinguishable or functionally equivalent from the perspective of the end user. For example, data intake and query

system **108** may be re-architected to run in a stateless, containerized environment. In some of these embodiments, data intake and query system **108** may be run in a computing cloud provided by a third party, or provided by the operator of the data intake and query system **108**. This type of cloud-based data intake and query system may have several benefits, including, but not limited to, lossless data ingestion, more robust disaster recovery, and faster or more efficient processing, searching, and indexing. A cloud-based data intake and query system as described in this section may provide separately scalable storage resources and compute resources, or separately scalable search and index resources. Additionally, the cloud-based data intake and query system may allow for applications to be developed on top of the data intake and query system, to extend or enhance functionality, through a gateway layer or one or more Application Programming Interfaces (APIs), which may provide customizable access control or targeted exposure to the workings of data intake and query system **108**.

[0177] In some embodiments, a cloud-based data intake and query system may include an intake system. Such an intake system can include, but is not limited to an intake buffer, such as Apache Kafka® or Amazon Kinesis®, or an extensible compute layer, such as Apache Spark™ or Apache Flink®. In some embodiments, the search function and the index function may be separated or containerized, so that search functions and index functions may run or scale independently. In some embodiments, data that is indexed may be stored in buckets, which may be stored in a persistent storage once certain bucket requirements have been met, and retrieved as needed for searching. In some embodiments, the search functions and index functions run in stateless containers, which may be coordinated by an orchestration platform. These containerized search and index functions may retrieve data needed to carry out searching and indexing from the buckets or various other services that may also run in containers, or within other components of the orchestration platform. In this manner, loss of a single container, or even multiple containers, does not result in data loss, because the data can be quickly recovered from the various services or components or the buckets in which the data is persisted.

[0178] In some embodiments, the cloud-based data intake and query system may implement tenant-based and user-based access control. In some embodiments, the cloud-based data intake and query system may implement an abstraction layer, through a gateway portal, an API, or some combination thereof, to control or limit access to the functionality of the cloud-based data intake and query system.

3. Asset Hierarchy Monitoring and Reporting System

[0179] Aspects of embodiments heretofore described may be advantageously implemented with subject matter next discussed to provide novel embodiments related to the monitoring and reporting for an asset hierarchy. Inventive subject matter will become manifest through the description and discussion of an asset hierarchy monitoring and reporting system. The disclosed system operates automatically to receive, collect, or ingest computer readable data as might be relevant to one or more assets in the asset hierarchy, and operates automatically to perform processes or methods that operate against the data to provide effective monitoring and reporting for the asset hierarchy. The disclosed system has its operation controlled by certain command, control, and

configuration (CCC) information, in computer storage. The disclosed system may implement a control console function to enable a system user or administrator to create, view, and edit CCC information as necessary to determine the operation of the asset hierarchy monitoring and reporting system (AMRS).

[0180] In one aspect of the disclosed embodiments, collections of data accessible via a DIQ are queried via a control console interface with results reported to the user. The user may interactively optimize the search and classify certain information in the results. Using the optimized query and information classifications, command console functions can proceed to automatically construct an asset hierarchy that, as part of the CCC information, may determine future asset hierarchy monitoring and reporting operations. This computer aided and implemented bottom-up construction of an asset hierarchy from volumes of data as may be extant in a DIQ system represents a vast improvement over an asset monitoring system implementation that requires manual determination of an asset hierarchy in a top-down approach, perhaps through laborious ETL (extract/transform/load) efforts, schema developments, and the like. Actual machine data reflecting the real-life asset hierarchy is harvested to derive the asset hierarchy definition for monitoring and reporting operations moving forward rather than requiring the administrative definition of asset hierarchy constructs to which relevant data can then be associated if the constructs are correct and the data conforms.

3.1 System Overview

[0181] FIG. 10 illustrates an asset hierarchy monitoring and reporting system deployment in one embodiment. Block diagram **2100** includes assets and asset data generators **2110**, secondary systems **2112**, intermediary data system **2114**, intermediary data store **2116**, data intake and query system (DIQ) **2120**, asset monitoring and reporting system (AMRS) **2140**, and user interface apparatus **2105**. DIQ **2120** is shown further to include event data store **2122**, metrics data store **2124**, command, control, and configuration (CCC) data store **2132** (independently or shared with AMRS **2140**), and CCC console processor **2134** (independently or shared with AMRS **2140**). AMRS **2140** is shown further to include command, control, and configuration (CCC) data store **2132** (independently or shared with DIQ **2120**), CCC console processor **2134** (independently or shared with DIQ **2120**), monitor/reporter processor **2142**, and asset hierarchy information **2150** shown to be included within its CCC data store **2132**.

[0182] Assets and asset data generators **2110** may include real-world physical assets for which monitoring and/or reporting are desired. Examples may include active and passive items, structures, machinery, components, devices, parts, assemblies, and interconnections therefore, as may be utilized for an oil drilling platform, a transmission pipeline, a factory or assembly line, an electrical generation facility (e.g., a windmill farm, hydroelectric plant, or nuclear generation facility), an electrical transmission facility or grid, a sensor deployment (e.g., seismographic stations, weather stations, oceanographic sensor buoys, GPS trackers), a satellite network, a chemical processing lab or facility, a distribution of Internet-of-Things (IoT) devices, a vehicle, a hospital, a region, or a person, to name but a few examples. An asset may generate data about itself or about another asset, implicitly or explicitly, as part of its principal opera-

tion and functioning, or as part of built-in monitoring and diagnostic capabilities. An asset may be a discrete asset, such as a temperature probe or other sensor, or may be a collection of other assets and/or other collections of assets, such as a distribution pipeline asset which is made up of multiple regional distribution segments, which are each made up of many pipes, valves, pumps, sensors, and so on, and the like. In terms of asset monitoring and reporting, an asset may also be an information asset, such as a single endpoint reading from a set of multiple different readings produced by a single sensor device. While many types of assets are possible, the assets and asset data generators 2110 in a particular embodiment, implementation, application, installation, or instance, may include the assets that commonly belong to a system or domain to be monitored, such as a nationwide common carrier fleet, one or more of a company's manufacturing plants, an early warning sensor network, or a municipality's sewer treatment and water filtration plant. Information produced by or about assets may be communicated directly or indirectly via some arrangement or combination of data polling, receiving, consolidating, concentrating, integrating, collecting, forwarding, filtering, and processing devices, and the like, such as a process control computer at a manufacturing plant, which may itself be an asset; and producing asset information by such devices may simply involve the presentation or transmission of asset information without creating or originally generating it. Regardless, the asset information appears at the intake side of the data intake and query system 2120. The arriving data for a system or domain may likely appear in different forms, in different patterns, at different times, with different timing factors, structured and unstructured (e.g., freeform, textual, variable length, etc.), crude and processed, simplex and complex, standards-conforming and not, via multiple communication means, modes, and methods, for ingestion by the DIQ system 2120. Additional asset-related information may be injected into the conceptual flow of data to the DIQ by secondary systems such as 2112. Secondary systems 2112 may include, for example, the service billing system of an outside contractor that reports repairs to assets of 2110. Secondary systems 2112 may or may not be linked electronically to the assets of 2110 but, nonetheless, may produce information related to those assets which may be ingested by DIQ 2120. Further, block diagram 2100 illustrates that information in the conceptual flow from the system/domain assets 2110 to DIQ 2120, may arrive at the DIQ directly or via an intermediary data system such as 2114. In one embodiment, intermediary data system 2114 may act, for example, as a pass-through, information aggregating device. In one embodiment, intermediary data system 2114 may act, for example, as a store-and-forward device that accumulates asset information in data store 2116 before forwarding it on to DIQ 2120 according to its operation and programming. In one embodiment, intermediary data system 2114 may be, for example, and OPC/UA server using standard protocols to collect asset information which it may later provide to DIQ 2120. Intermediary data systems such as 2114 may be relatively dumb devices in the sense of performing only rudimentary operations on their input data, sending it back out largely unchanged from how it was received. Intermediary data systems such as 2114 may be relatively smart devices in the sense of performing more advanced operations on their input data, sending it back out

in a possibly highly processed form from that which was received. Many embodiments are possible.

[0183] Whether originating from assets or asset data generators of 2110, or from secondary systems 2112, and whether arriving directly or via an intermediary system such as 2114, asset-related data arrives at data intake and query 2120 for intake, ingestion, searching, monitoring, reporting, and/or other processing. In an embodiment, DIQ system 2120 may well be an implementation of a data intake and query system such as system 108 as shown and discussed in relation to FIGS. 1, 2, and 9, or system 1006 as shown and discussed in relation to FIG. 8. In one embodiment, where DIQ system 2120 is implemented after the fashion of DIQ system 108 of FIG. 9, asset data for 2110 of FIG. 10 may include one or more data sources (202 of FIG. 9) coupled to one or more forwarders (204 of FIG. 9). In such an embodiment, asset data for 2110 of FIG. 10 may include one or more data sources, perhaps intermediary system 2114, coupled to one or more ERP processes (e.g. 1110, 1112 of FIG. 9). Many arrangements and combinations are, of course, possible in light of the disclosure herein.

[0184] In an embodiment, asset data received for intake by DIQ 2120 of FIG. 10 may arrive at event data store 2122, which may be implemented as an indexed data store such as data store 208 of FIG. 2. In an embodiment, asset data received for intake by DIQ 2120 of FIG. 10 may arrive at a metrics data store 2124. Metrics data store 2124 may be implemented by DIQ 2120 as an alternative or supplementary data storage regime to event data store 2122 and may have superior performance aspects when processing large volumes of, or data chiefly consisting of, metrics data. The operation of DIQ 2120 to perform intake and subsequent processing of the received asset data is controlled by information in command/control/configuration data store 2132. CCC console 2134 effectively implements the control panel for the system by providing user interfaces and related processing that enable user, such as a system administrator or operator, to view, create, edit, delete, or otherwise process or manipulate the information in this CCC data store 2132 that controls the operation of the system. CCC console 2134 may cause the display of a user interface on a dedicated console device or on a multi use device such as a network attached user computer as depicted by 2105 of FIG. 10. CCC data store 2132 and console processor 2134 may be shared in whole or part by DIQ 2120 and asset system 2140, as depicted in FIG. 10 by the straddling of the boundary 2130 between the two. Shared boundary 2130 illustrates the coupling between the functionality described for each of DIQ 2120 and asset system 2140. While DIQ 2120 and asset system 2140 are depicted and described as distinct components in block diagram 2100 one of skill can appreciate that implementations may vary. In one embodiment, DIQ 2120 and asset system 2140 are distinct systems, running on distinct platforms, that do not share CCC data store 2132 or CCC console processor 2134, and that are coupled by defined interfaces communicating over generalized network facilities. In one embodiment, DIQ 2120 and asset system 2140 are completely integrated and installed or formed together as a single unit. In one embodiment, software of DIQ 2120 may be installed separately and in advance of asset system 2140 software, and asset system software may be packaged for installation as an application or subsystem of DIQ 2120. Accordingly, distinctions illustrated and described here represent logical distinctions that are useful

to convey an understanding of inventive aspects but do not impose limitations on the practice of those aspects.

[0185] Like data intake and query system 2120, the operation of asset system 2140 is directed and determined by certain information in a command, control, and configuration data store such as 2132. If asset system 2140 does not share CCC data store 2132 with DIQ 2120, it may have its own. Similarly, if asset system 2140 does not share console processor 2134 with DIQ 2120, it may have its own, which may nonetheless utilize user interface device 2105 for displaying interactive user interfaces. Information of CCC data store 2132 for asset system 2140 is shown to include information defining and representing an asset hierarchy 2150. Asset hierarchy definition 2150 is a logical construct that may be implemented as one or more data structures, having one or more formats, stored at one or more locations, across one or more device types. Sufficient information exists in the stored representation of asset hierarchy definition 2150, directly or indirectly, expressed or implied, to enable a computing machine arrangement to operate in accordance with the logical asset hierarchy definition 2150.

[0186] Monitor/reporter 2142 of asset system 2140 utilizes asset hierarchy definition 2150 and perhaps other information in CCC data store 2132 to direct, determine, condition, or otherwise influence its operation to effect monitoring and/or reporting related to one or more assets included in an asset hierarchy. The activity of monitor/reporter 2142 may be variously performed as continuous, intermittent, scheduled, or on-demand processing. Such monitoring and/or reporting activity may produce outputs immediately intended for human consumption, such as a status display user interface presented on a user interface device such as 2105, or outputs immediately intended for machine use, such as an event record recorded to the event data 2122 of DIQ 2120 in response to detecting a condition during ongoing analysis of incoming asset data. Monitor/reporter 2142 may perform any of its functional processing directly or may interface with other systems and/or subsystems, such as DIQ 2120, to have certain functional processing performed. These and other embodiments are possible.

[0187] In an embodiment, DIQ 2120 and asset system 2140 may be implemented as dedicated hardware, dedicated computing hardware programmed with software, general purpose and/or mixed-use computing hardware specialized with software to implement operation as a DIQ and asset system, or the like, alone or in combination.

[0188] FIGS. 11A and 11B illustrate an illustrative asset hierarchy structure. FIGS. 11A and 11B depict the same example asset hierarchy structure, that is to say, the same set of nodes with the same set of interrelationships. FIG. 11A, however, labels each node according to its role while FIG. 11 B labels each node according to one illustrative physical world example. Hierarchical node 2202 is designated as the system or root node and is the only node of hierarchy 2200 that does not have a parent node (does not descend from a superior node). Root node 2202 may or may not represent a unitary physical asset, and often may not, but in any event may represent the aggregation or collection of all that is represented in the hierarchy of nodes descending from it. In an embodiment, every asset tree or hierarchy may have a system or root node, and the system or root node may serve as a principal entry point for navigating the hierarchy and for identifying the hierarchy. The system or root node may represent the system, domain, physical asset, category, loca-

tion, owner, or other characteristic, designation, or construct, that may unify all that is represented in the hierarchy of nodes descending from it. In an embodiment, the root node of an asset hierarchy may represent the system against which monitoring and/or reporting may be conducted, such as a water filtration system. In an embodiment, the root node of an asset hierarchy may represent the domain of assets against which monitoring and/or reporting may be conducted, such as the Western region air-quality sensor stations. System/root node 2202 is shown to have two child (immediate descendent) nodes, 2210 and 2212, each of which has its own children. Nodes in hierarchy 2200 that have one or more child nodes are identified in FIG. 11A as asset/container nodes. (Systems/root node 2202 fits the criteria and may be considered an asset/container node though not labeled as such.) An asset/container node may represent a physical construct having subordinate assets such as a machine that has subassemblies and/or parts. An asset/container node may represent a logical construct having subordinate assets such as a category, class, or type of assets, such as “pumps” or “filters,” and serves as a container or collection point for such assets. Asset/container node 2210 is shown to have two child nodes, 2230 and 2232, neither of which has any child node. Nodes in the asset hierarchy or tree that have no children are designated as leaf or terminal asset nodes. While a leaf or terminal asset node in the hierarchy may represent, for example, a physical object for which subordinate items, parts, sections, portions, subassemblies, data item/channels/streams, or the like, may indeed be capable of being identified, those subordinate elements may be effectively disregarded for monitoring and/or reporting purposes by the asset system because of their exclusion from the asset tree. Asset hierarchy 2200 is further shown to include three child nodes of asset/container node 2212: asset/container node 2234, and terminal asset nodes 2236 and 2238. Asset/container node 2234 is further shown to have child terminal asset node 2240.

[0189] Asset hierarchy 2200 may also be considered in terms of levels within its hierarchy. According to one paradigm, the levels of the hierarchy may be designated according to the number of steps a node is distanced from the root note. (A step may be considered to be the traversal of an internodal edge of the hierarchical graph.) System/root node 2202 of asset hierarchy 2200 occupies level 0 as it has no distance from itself. The direct children of system/root node 2202, i.e. nodes 2210 and 2212, may be considered level 1 nodes as each is distanced one step from system/root node 2202. Each of the direct children of the level 1 nodes (the “grandchildren” of systems/root node 2202) may be considered a level 2 node as each is distanced two steps from system/root node 2202. Finally, in this example, leaf node 2240 may be considered a level 3 note as it is the direct child of a level 2 node (2234) and is distanced three steps from system/root node 2202 (a “great-grandchild”).

[0190] FIG. 11B duplicates asset hierarchy structure 2200 from FIG. 11A, but now labeling each node according to a simplistic example to aid in understanding. In this example asset tree 2200 of FIG. 11B will represent a vacuum packing machine and, accordingly, root node 2202 is identified as “Vacuum Packer.” The vacuum packer is shown to include immediate component assets, Chamber Sensor 2210 and Pump Ass[embly] 2212. Chamber Sensor 2210 is shown to include Temp[erature] and Pressure assets, 2230 and 2232, each of which may represent a source, stream, channel, or

component of measurement data for temperature and pressure metrics, respectively, produced by the chamber sensor device represented by node 2210. Pump assembly 2212 is shown to include immediate component assets motor 2234, piston cylinder 2236, and valve 2238, each of which in this example is suggested to represent a physical entity. Node 2234, and by representation the motor of the pump assembly of the vacuum packer, is shown to include immediate component asset "Thermo [switch]" 2240 which may represent an overheat switch (or its state/signal) that is built into the motor represented by node 2234.

[0191] FIG. 12 illustrates methods of an asset hierarchy monitoring and reporting system in one embodiment. The illustrated methods may represent a significant portion of the processing workflow to establish and operate an asset monitoring and reporting system (AMRS). Flowchart 2300 of FIG. 12 may beneficially be discussed in terms of an AMRS after the fashion illustrated in FIG. 10, such as an AMRS making combined use of functioning described for the data intake and query system 2120 and the asset system 2140 of FIG. 10.

[0192] At block 2320 of FIG. 12, the sources of asset-related data are configured so that DIQ functions can intake and process the data. The processing of block 2320 may include storing relevant information in command, control, and configuration (CCC0 data store 2322. In an embodiment, such information may include data source definitions, data model information, field mappings and extraction rules, or the like, as may be useful to control the operation of the DIQ to intake and process, for storage and query, asset-related data from any number and variety of data sources. The processing of block 2320 may include causing the display of user interfaces on a user interface device, such as user computer 2312. The displayed user interfaces may be interactive, enabling a user to provide inputs to the processing of block 2320. In an embodiment user inputs may be content or indicators for data items, selections, commands, and such. The processing of block 2320 may utilize certain user inputs in the immediate context to condition, influence, or direct CCC console functions performed at block 2320, and may utilize the same or other certain user inputs to modify information in the CCC data store and thereby affect AMRS operation beyond the immediate context. As suggested, the method processing of block 2320 may be performed in an embodiment by a CCC console processor such as 2134 of FIG. 10. Block 2320 of FIG. 12 may utilize a user interface component such as user interface display 2400 of FIG. 13 in the course of its processing.

[0193] FIG. 13 illustrates a user interface display of a console function for specifying data inputs. Interface display 2400 is shown to include system header bar 2402 and header area 2410, followed by a detail display area. The system header bar 2402 is shown to include a system name, "splunk>", an application/function drop-down selection element 2404, a system menu area 2406, and a system search criteria entry box 2408. Header area 2410 is shown to include the title, "Data inputs", for the user interface. The detail display area is shown to include section headers 2412 and 2414 which may be used to identify and delineate different sections, portions, panels, or the like, of the detail display area. Section header 2412 identifies a "Local inputs" section of the detail display area which may include a tabular display identifying the types of local input data sources that may be defined to a DIQ in one embodiment. The tabular

display of user interface 2400 includes column header row 2418 identifying "Type", "Inputs", and "Actions" columns. The "Actions" column for each row contains an "Add new" interactive element that enables a user to signal her desire to define a new data source/input for the DIQ of the type identified in the row. User interaction with the "Add new" interactive element may cause the display of, or navigation to, a user interface display component that enables a user to interact with the system for the addition of information to the CCC data store to effect processing for a new data input of the type identified in the row. The "Inputs" columns for each row contains a count of the number of already-defined data sources/inputs for the DIQ of the type identified in the row. At row 2420, the "Type" column includes an interactive identifier, "Files & directories", referring to data inputs/sources that are files, or directories of files, in the file system of one or more host computers. In one embodiment, user interaction with a type identifier, such as "Files & directories", may result in the appearance of, or navigation to, a user interface display component that directly provides, or provides access to, one or more user interfaces for viewing, editing, deleting, or otherwise interacting with information of a CCC data store related to already-defined data sources/inputs of the type identified in the row. At row 2422, the "Type" column includes an interactive identifier, "HTTP Event Collector", referring to data input/sources that are received over HTTP or HTTPS connections. At row 2424, the "Type" column includes an interactive identifier, "TCP", referring to data input/sources that are received over a listened-to TCP port. At row 2426, the "Type" column includes an interactive identifier, "UDP", referring to data input/sources that are received over a listened-to UDP port. At row 2428, the "Type" column includes an interactive identifier, "Scripts", referring to data input/sources that are executions of scripts or other programming that collect or generate data. At row 2430, the "Type" column includes an interactive identifier, "OPC UA Pull Connect", referring to data input/sources that are collected from an OPC/UA server. At row 2432, the "Type" column includes an interactive identifier, "OPC UA Event Notification", referring to data input/sources that are event notification received from an OPC/UA server. At row 2434, the "Type" column includes an interactive identifier, "OPC UA Subscription", referring to data input/sources that are subscribed from an OPC/UA server. At row 2436, the "Type" column includes an interactive identifier, "UA Simulator Server", referring to data input/sources that are received by an OPC/UA server simulation.

[0194] Without indicating any special significance or importance, it is noted that the OPC/UA data input types designated in rows 2430, 2432, 2434 and 2436, may be relevant to a large class of use cases for an asset monitoring and reporting system (AMRS) as described herein, as the types of installations and systems that utilize OPC/UA standards and protocols are the types of installations and systems having numbers of assets for which data is generated and for which robust monitoring and/or reporting can be beneficial. One large class of installations and/or systems having numbers of assets for which data is generated and for which robust monitoring and/or reporting can be beneficial, whether utilizing OPC/UA or not, are industrial installations and related industrial control systems. Such systems may include remote monitoring and control (M & C) systems designed to control large or complex facilities such as

factories, power plants, network operations centers, airports, and spacecraft, with some degree of automation. Such a remote monitoring and control system may include a supervisory control and data acquisition (SCADA) system that operates with coded signals over communication channels to perhaps acquire information about the status of remote equipment and perhaps to issue coded command signals to remote equipment, possibly over large distances, and possibly using protocols and devices supporting OPC/UA. OPC/UA is well understood in the art as an industrial machine-to-machine (M 2M) facility, and information regarding OPC/UA is promulgated by various means including the website found on the internet at the opcfoundation.org domain operated by the OPC Foundation, headquartered at 16101 N. 82nd Street, in Scottsdale, Arizona.

[0195] It is understood that in their role and to the extent they concern assets of an asset hierarchy, industrial control systems, including remote monitoring and control systems and SCADA systems, may perform their own monitoring and reporting regarding assets in an asset hierarchy. Such monitoring and reporting is distinct from the monitoring and reporting performed by an AMRS implementation as described and taught here, even while the one may provide data to the other.

[0196] At block 2330 of FIG. 12, data sources with asset-related data are processed by DIQ functions. The processing of block 2330 may access information of CCC data store 2322 to condition, control, direct, configure, or otherwise influence the intake and processing of the data from the data sources. Access to information of CCC data store 2322 may include access to information added or modified during the processing of block 2320. In an embodiment, processing of block 2330 may be initiated automatically (i.e., without user intervention) and may be performed in a continuous, periodic, regular, scheduled, timed, intermittent, or other such fashion automatically. In an embodiment, the processing of block 2330 may be performed on an on-demand basis. Processing of block 2330 may result in the reflecting of asset machine data in computer storage of any number and variety of types. The processing of block 2230 may be performed by a DIQ, such as 2120 of FIG. 10. The asset machine data 2332 of FIG. 12 may be stored as event data 2122 and/or as metrics data 2124 of DIQ 2120 of FIG. 10, for example, in an embodiment.

[0197] At block 2340 of FIG. 12, an asset tree structure may be created or edited. If an asset tree structure is being created, block 2340 may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such, for the creation of a new asset tree by the processing of block 2340. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If an asset tree structure is being edited, information of an existing asset tree in CCC data store 2322 may be directly or indirectly represented in a display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 2340 to construct a proper representation of a new or updated asset tree, and a representation of the new or updated asset tree may be reflected in the information of CCC data store 2322. User interface displays of the processing of block 2340 may be caused to be displayed on a user interface device such as user computer 2314, for example. In an embodiment user

computer 2312 and user computer 2314 may be the same device or two or more different devices. Embodiments may vary as to the minimum and total amount of information that may be included for each node of an asset tree. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for an asset tree representation, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 2340, or elsewhere and at other times. The processing of block 2340 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as shown in FIG. 10. At the conclusion of the processing of block 2340 of FIG. 12, in an embodiment, information representing a new, changed, or unchanged asset tree definition may be found in CCC data store 2322.

[0198] At block 2342 of FIG. 12, definitions for asset metrics may be created or edited. If one or more metric definitions are being created, block 2342 may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation of new metrics definitions by the processing of block 2342. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If the definition of one or more metrics is being edited, information of the existing one or more metrics in CCC data store 2322 may be directly or indirectly represented in a display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 2342 to construct a proper representation of one or more new or updated metric definitions, and a representation of the new and/or updated metric definitions may be reflected in the information of CCC data store 2322. User interface displays of the processing of block 2342 may be caused to be displayed on a user interface device such as user computer 2314, for example. Embodiments may vary as to the minimum and total amount of information that may be included for each metric definition. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of a metric definition, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 2342, or elsewhere and at other times. The processing of block 2342 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as shown in FIG. 10. At the conclusion of the processing of block 2342 of FIG. 12, in an embodiment, information representing one or more new, changed, and/or unchanged metric definitions may be found in CCC data store 2322.

[0199] In an embodiment, the processing of block 2342 may include processing to create or edit definitions, specifications, indications, or the like, for one or more associations between and among one or more metrics, on the one hand, and one or more assets represented in an asset hierarchy, on the other hand. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like, that are or may be used for the representation of such associations in computer storage at any place and time. In an embodiment, defined associations are reflected in the information of CCC data store 2322, and here, as elsewhere, the representation of the

defined associations in computer storage may be direct or indirect, expressed or implied, or otherwise.

[0200] At block 2344 of FIG. 12, definitions for conditions and/or alerts and/or actions may be created or edited. If one or more conditions and/or alerts and/or actions are being created, block 2344 may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation by the processing of block 2344 of new condition and/or alert and/or action definitions. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If the definition of one or more conditions and/or alerts and/or actions is being edited, information of the existing one or more conditions/alerts/actions in CCC data store 2322 may be directly or indirectly represented in display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 2344 to construct a proper representation of the one or more new or updated condition/alert/action definitions, and a representation of the new and/or updated condition/alert/action definitions may be reflected in the information of CCC data store 2322. User interface displays of the processing of block 2344 may be caused to be displayed on a user interface device such as user computer 2314, for example. Embodiments may vary as to the minimum and total amount of information that may be included for each condition/alert/action definition. Embodiments may also vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of a condition/alert/action definition, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 2344, or elsewhere and at other times. The processing of block 2344 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as shown in FIG. 10. At the conclusion of the processing of block 2344 of FIG. 12, in an embodiment, information representing one or more new, changed, and/or unchanged condition/alert/action definitions may be found in CCC data store 2322.

[0201] In an embodiment, the processing of block 2344 may include processing to create or edit definitions, specifications, indications, or the like, for one or more associations between and among one or more conditions/alerts/actions, on the one hand, and one or more other defined objects, elements, or constructs, on the other hand. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like, that are or may be used for the representation of such associations in computer storage at any place and time. In an embodiment, defined associations are reflected in the information of CCC data store 2322, and here as elsewhere, the representation of the defined associations in computer storage may be direct or indirect, expressed or implied, and otherwise.

[0202] At block 2346 of FIG. 12, command, control, and configuration (CCC) information for the monitoring and/or reporting processing of the AMRS as may appear in CCC data store 2322, may be created or edited. If the information is being created, block 2346 may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation of new CCC information by the processing of block

2346. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If CCC information is being edited, existing information of CCC data store 2322 may be directly or indirectly represented in a display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 2346 to construct a proper representation of the information and to reflect it in CCC data store 2322. User interface displays of the processing of block 2346 may be caused to be displayed on a user interface device such as user computer 2314, for example. Embodiments may vary as to the minimum and total amount of user-configurable information, and the purposes of which, that may be included in CCC data that determines, conditions, or otherwise influences the operation of monitoring and/or reporting aspects of the processing of an AMRS. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of monitoring and/or reporting CCC information, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 2346, or elsewhere and at other times. The processing of block 2346 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as shown in FIG. 10. At the conclusion of the processing of block 2346 of FIG. 12, in an embodiment, information representing new, changed, and/or unchanged CCC information that determines, conditions, or otherwise influences the operation of monitoring and/or reporting aspects of the processing of an AMRS may be found in CCC data store 2322.

[0203] It is noted that as the illustrative example 2300 of FIG. 12 envisions an AMRS combining DIQ and asset system functionality as described in relation to FIG. 10, CCC data store 2322 of FIG. 12 finds a counterpart in a shared CCC data store 2132 as described in relation to FIG. 10, and, similarly, any discussion here in reference to CCC console processor 2134 of FIG. 10 embraces an embodiment of a shared CCC console as discussed there.

[0204] At block 2350 of FIG. 12, ongoing automatic and/or on-demand monitoring and/or reporting for an asset hierarchy as may be provided by an AMRS is conducted. The processing of block 2350 may utilize information reflected in CCC data store 2322 as described for the processing of other blocks of 2300, to determine, direct, condition, or otherwise influence its operational activity. In an embodiment, asset hierarchy monitoring activity may include data intake, internal data generation, computer-to-computer data transmission/presentation, and computer-to-person data presentation, for example, with a possible emphasis on data intake and generation aspects. In an embodiment, asset hierarchy reporting activity may include one or more of the same with a possible emphasis on data presentation aspects. The processing of block 2350 may be conducted, for example, by a combination of DIQ 2120 and asset system 2140 of FIG. 10, in an embodiment.

[0205] Examples for the processing of block 2350 for illustrative implementations, applications, or instantiations of an AMRS, follow. In one example, the processing of block 2350 may include monitoring sensors on a connected soldier in the battlefield and generate appropriate alerts for received impacts or heartbeat abnormalities. In one example, the processing of block 2350 may include monitoring pro-

cess control equipment in a beverage processing facility and generate alerts for out-of-range temperatures or vibration anomalies. In one example, the processing of block 2350 may include monitoring activity in a workorder management system report impacts caused by scheduled downtime. In one example, the processing of block 2350 may include monitoring changes in process control settings, such as PID control settings, and report impacts on the accuracy of control loops. These are but a few illustrative examples.

[0206] An appreciation for the methods 2300 of FIG. 12 may be further developed by consideration of methods and user interface examples illustrated and discussed in relation to figures that follow.

3.2 Asset Hierarchy Establishment

[0207] FIG. 14 illustrates a method for constructing an asset tree representation in control storage. Method 2500 of FIG. 14 illustrates a method as might be employed in an embodiment during the processing of block 2340 of FIG. 12, for example. Processing of block 2510 of FIG. 14 causes the display of a user interface on a user interface device. The user interface may be interactive enabling a user to both receive information from the AMRS (e.g., its CCC console processor) and to provide information to the AMRS as enabled by the user interface. At block 2512, the AMRS receives an indication of search criteria which may be provided by user interaction with the user interface. In an embodiment, the search criteria indications may represent some or all of the search criteria used in a DIQ search query to identify useful asset-related information. In an embodiment, the search criteria indications may be provided by user interaction in one or more forms, for example, selections of checkboxes associated with a particular search criteria, segments of text in the form of a query language that specify search criteria, the text of a fully formed search query in accordance with the query language, or others. At block 2514, the indications of search criteria received at block 2512 are utilized by the AMRS to construct, as necessary, a search query to locate useful asset-related information, and to execute a search to locate such useful asset-related information in accordance with the indicated search criteria. In an embodiment, the AMRS may conduct such a definition-time search by passing a search query request to a DIQ component of the system. At block 2516, some or all of the search results from the processing of block 2514 are caused to be displayed via a user interface. In an embodiment, processing of block 2516 enables a user to view the search results and to iterate back to block 2512 if the search results are deemed unsatisfactory. Processing of block 2516 may enable a user to indicate satisfaction with the search results and thereby advance to processing that enables a user to make certain classifications of the search results, perhaps by updating the display of the user interface. In one embodiment, the user interface initially displayed from the processing of block 2516 may enable the user to immediately indicate the certain classifications, and any particular user interactions to indicate the classifications may be interpreted by block 2516 as an indication by the user of satisfaction with the search results. Other embodiments are possible. At block 2518, indications of user classifications of the search results are received and processed. The processing of block 2518, in an embodiment, may receive and process an indication of data in the search result set that may be used to provide an identification of assets to be represented as

nodes in an asset hierarchy. The processing of block 2518 in an embodiment may receive and process an indication of data in the search result set that may be used to provide an identification of a parent asset corresponding to an asset identified in the search result set. In an embodiment, the user indications received and processed at block 2518 may be indications of the identification of fields produced by the search query that provide asset identifying and parent asset identifying information. At block 2520, the AMRS makes a determination of an asset tree structure by processing the result of the search query produced at block 2514, or a related search, in view and consideration of the asset and parent indications received at block 2518. A related search, in an embodiment, may be a search that is derived from the original search but is somehow modified, perhaps by restricting the fields returned in the search result set, or perhaps by expanding the scope of the data searched, or perhaps by other variations. The processing of block 2520, in an embodiment, may determine all of the unique identifiers that may be found within the designated asset identifier field of the search result set and create an asset hierarchy node for each, determine a respective parent for each unique asset using information found within the designated asset parent identifier field of the search result set, and cross reference asset identifiers and asset parent identifiers to determine the hierarchical relationships for the asset nodes and create a representation of those associations between the created asset nodes. Such representations may be express or implied, direct or indirect, or otherwise. In one embodiment, such representation is made by including the asset id for the parent asset among the information of each non-root node. In one embodiment, such representation is made by including the asset id for the child among one or more entries of a child-list maintained for each asset node. Other embodiments are possible. The processing of block 2520, in an embodiment, may conclude with a representation of the determined asset hierarchy recorded in computer storage, perhaps in computer storage of the local working context, or perhaps in persistent computer storage of a CCC data store. These and other embodiments are possible.

[0208] At block 2522, in an embodiment, processing is performed to receive and process user indications of data items, elements, fields, constants, or the like that should be included in, directly or indirectly, the definitional information of nodes in the asset hierarchy tree. In one embodiment, indications received and processed at block 2522 are indications of fields in the search result set produced at block 2514 or produced elsewhere that contain information that should be included among the definitional information of the asset hierarchy. The indications received and processed at block 2522 are utilized in the processing of block 2524 to augment the asset identifier information of the asset hierarchy with additional, possibly per-node, information. The processing of block 2524, in an embodiment, may conclude with a representation in the computer storage that reflects an asset hierarchy structure with augmented information. These and other embodiments are possible.

[0209] FIG. 15 illustrates a user interface display for an asset search console function. User interface display 2600 is such as might be caused to display during the processing of block 2340 of FIG. 12, or the processing of method 2500 of FIG. 14, for example; such processing possibly performed by a command, control, and configuration console processor such as 2134 of FIG. 10. User interface display 2600 of FIG.

15 is shown to include system header bar **2402**, application information and menu bar **2602**, console function header area **2610**, search specification area **2620**, and search result area **2630**. Application information and menu bar **2602** is shown to include a title for the application or topic within the system to which the user interface pertains (“Asset Intelligence”), and a number of menu and/or navigation options (“Configure”, “Monitor”, “Diagnose”, “Search”) which may be interactive elements for navigation, drop-down menus, drop-down selection lists or such. Console function header area **2610**, which may identify an AMRS CCC console function to which the current user interface pertains, is shown to include function identifier or title **2612** (“Asset import”), function/task progress indicator **2614**, and “Next” action button **2616**. Function/task progress indicator **2614** may provide a user with information about the association of the current user interface **2600** with a particular step, segment, or subtask of a multipart task, process, or workflow, such as “Asset import.” As shown, progress indicator **2614** indicates by a solid colored circle that interface **2600** is associated with a “Search assets” portion of an “Asset import” function, task, or workflow. In an embodiment, action button **2616** may be interactive enabling the user by a keypress, mouseclick, touchscreen press, or the like, to indicate to the AMRS a desire to navigate to processing and an associated user interface for a subsequent portion of the “Asset import” function.

[0210] Search specification area **2620** is shown in FIG. 15 to include a search specification text box **2622** displaying the text of a search language query **2624**, a search timeframe specification component **2626**, and an execute-search action button **2628**. In an embodiment, the displayed text of the search language query **2624** (or other form of search criteria specification) may be a system supplied initial value, a user entered value, a user edited version of a system supplied value, a recalled value from a user profile, a last-used value, or a value from another source. Other implementations are possible. Search timeframe specification component **2626** may be an interactive element such as a drop-down selection box that displays some defaulted value or the last value specified or selected by the user, shown here as “Last 4 hours.” Execute-search action button **2628** may be an interactive element such as an iconized action button that may enable the user by a supported interaction to indicate to the AMRS a desire to execute the search specified by information that may include values associated with user interface elements **2622** and **2626**. An AMRS in an embodiment may respond to user interaction with action button **2628** by receiving user input and performing processing to execute a search specified at least in part by content represented in search specification area **2620** of interface **2600**. In an embodiment, performing a specified search for asset-related data may include invoking services or functions of a data intake and query system (DIQ) to perform the specified search and produce a search result. An embodiment may utilize the search result set to populate the search result display area **2630**. While inventive aspects may not be so limited, search result display area **2630** of example interface **2600** is shown to present search results in a tabular format, shown here as an orthogonal grid of information cells each uniquely identifiable by the combination of the row in which it appears and the column in which it appears. Search result display area **2630** is shown to include column headings area **2632**, table data area **2634**, and table data navigation control

area **2636**. Each of the rows **2650a-j** appearing in table data area **2634** has an information display cell located in each of columns **2640a-h**, indicated by the headings of **2632** to be information cells for “Equipment”, “Equipment Category”, “Manufacturer”, “Model”, “Installation date”, “Flow”, “ORP_Level”, and “Pressure”, respectively. In an embodiment, the column headings of **2632** may be populated with the names of fields returned in the search result set. In an embodiment, each of the rows appearing in the table data area may correspond to an individual instance, record, entry, or such, of a search result. In an embodiment, individual column headings of **2632** may be interactive so as to enable the user to indicate to the AMRS a desire to see the information appearing in table data area **2634** sorted according to the order of the values appearing in the column corresponding to the heading with which the user interacted. These and other embodiments and variations are possible.

[0211] It may be worth noting that one of skill may well consider the present teachings, illustrated here in terms of a tabular information format with rows and columns, in different terms that may more artfully pertain in a particular implementation without necessarily departing from inventive aspects. For example, a table of data may be considered in terms of a file, dataset, array, list, collection, or others, and perhaps in particular reference to having one or more rows or row equivalents, and one or more columns or column equivalents. For example, a row may be considered in terms of the tuple, record, entry, line, dimension, or others. For example, a column may be considered in terms of a field, item, position, offset, dimension item, category, or others. For example, a column title or heading may be considered in terms of a field name, item name, or key. For example, a table cell may be considered in terms of a field value, item value, dimension item value, or value as may pertain to a key-value pair, or others. The illustration of inventive embodiments using a tabular data format with columns and rows should not be considered as limiting the practice of inventive aspects where that is not otherwise required.

[0212] A user may indicate to the AMRS that the search specified at **2620** is adequate to capture data that embraces an asset hierarchy for which the user desires monitoring and/or reporting operations to be performed by the AMRS. In the example illustrated by interface **2600**, a user may interact with action button **2616** to make such an indication of the adequacy of the specified search. In an embodiment, consequent to receiving an indication of user interaction with action button **2616** of interface **2600**, an AMRS may perform processing that directly or indirectly may cause the display of a user interface such as depicted in FIG. 16.

[0213] FIG. 16 illustrates a user interface display for an asset information classification console function. User interface display **2700** is such as might be caused to display during the processing of block **2340** of FIG. 12, or the processing of method **2500** of FIG. 14, for example; such processing possibly performed by a command, control, and configuration console processor such as **2134** of FIG. 10. User interface display **2700** of FIG. 16 is shown to include system header bar **2402**, application information and menu bar **2602**, console function header area **2610**, and search result area **2730**. System header bar **2402** and application information and menu bar **2602** are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Console function header area **2610** is as described

for the identically numbered element depicted in interface display 2600 of FIG. 15 with the exceptions that the current position indicator associated with function/task progress indicator 2614 is shown here advanced to the “Select columns” task portion, and that a “Previous” action button 2718 counterpart to “Next” action button 2616 now appears.

[0214] Search result display area 2730 of FIG. 16 largely corresponds to search result display area 2630 of FIG. 15. The column or field names appearing in column headings area 2632 of FIG. 15 also appear in the same relative positions in column names area 2732 of FIG. 16, albeit without the companion sort action interface elements. Table data area 2734 of FIG. 16 finds correspondence to table data area 2634 of FIG. 15 with a difference that table data area 2734 of FIG. 16 is shown to include a greater number of individual rows (2750 a-m). Columns 2740a-h of search result display area 2730 of FIG. 16 directly correspond to columns 2640a-h of FIG. 15. Notably, the column header area 2731 of FIG. 16 is expanded beyond merely including a column or field name header row 2732 (or 2632 of FIG. 15) to also including a column classification row 2733. Column classification row 2733 includes a classification interface component in each of the columns of the row, such as classification selection drop-down 2733a appearing in row 2733 at column position 2740a. User interaction with classification selection drop-down 2733a may cause the display of a drop-down selection list such as 2760. The drop-down selection list may provide a user with a number of interactive selection options, each option designating a role that column data may play in constructing a representation of an asset hierarchy. Drop-down selection list 2760 of the illustrated embodiment is shown to include a defaulted “Skip Column” option 2762a, the default to, or express selection of which, indicates that data of the column may be ignored for purposes of defining an asset hierarchy.

[0215] Drop-down selection list 2760 of the illustrated embodiment is shown to include a “Parent Asset” option 2762b, the selection of which indicates that data of the column includes an asset identifier for an asset that is the parent of the asset represented by the particular row in the table. When creating a representation of an asset hierarchy from the search result data, the AMRS may utilize information in a “Parent Asset” column when determining the internodal associations or relationships that define the hierarchical structure. In one embodiment, no more than one column may be classified as a “Parent Asset” column.

[0216] Drop-down selection list 2760 of the illustrated embodiment is shown to include an “Asset” option 2762c, the selection of which indicates that data of the column includes the asset identifier for the asset represented by the particular row in the table. When creating a representation of an asset hierarchy from the search result data, the AMRS may utilize information in an “Asset” column when determining or associating an identification for nodes in the hierarchy. In one embodiment, no more than one column may be classified as an “Asset” column. In one embodiment, a first column designated as an “Asset” column is used to determine or associate an identification for a node, and any subsequent column designated as an “Asset” column is included in the asset tree definitional data as nickname metadata for the node.

[0217] Drop-down selection list 2760 of the illustrated embodiment is shown to include a “Metric” option 2762d, the selection of which indicates that data of the column is

related to a metric for the asset represented by the particular row in the table. When creating a representation of an asset hierarchy from the search result data, the AMRS may utilize information in a “Metric” column, particularly the column name or heading in one embodiment, to include the metric among one or more metrics that may be associated with the asset represented by the particular row in the table. Such an association will not occur in an embodiment where the metric column value in the asset row is, for example, a null value.

[0218] Drop-down selection list 2760 of the illustrated embodiment is shown to include an “Asset Metadata” option 2762e, the selection of which indicates that data of the column includes information that describes or is otherwise related to the asset represented by the particular row in the table. When creating a representation of an asset hierarchy from the search result data, the AMRS may utilize information in an “Asset Metadata” column to populate a metadata portion, and possibly a user-defined metadata portion, of definitional data associated with a node of the asset hierarchy.

[0219] Classification interface components appearing in each of the columns of row 2733 operate likewise in an embodiment for their respective columns. In one embodiment, “Next” action button 2616 may be disabled for user interaction unless a column has been classified as an “Asset” column and a column has been classified as a “Parent Asset” column. When system requirements are satisfied, “Next” action button 2616 may be enabled. When user desires are satisfied, a user may interact with “Next” action button 2616 to indicate satisfaction. In one embodiment, and AMRS receiving an indication of user interaction with “Next” action button 2616 may conclude processing as described in relation to block 2518 of FIG. 14, and proceed to processing described in relation to block 2520 of FIG. 14 where the AMRS determines an asset tree hierarchy using search results and user classifications, in one embodiment.

[0220] In an embodiment employing user interface 2700 of FIG. 16, the AMRS may perform processing as described in relation to blocks 2522 and 2524 of FIG. 14 in conjunction with the processing as described for blocks 2518 and 2520. In one embodiment, for example, a user input indicating a classification of the column as a “Metric” or “Asset Metadata” column may implicate the processing of block 2522 of FIG. 14. Other implementations and embodiments are possible. In one embodiment, after determining an asset hierarchy from search result data the AMRS may cause the display of the user interface as illustrated in FIG. 17.

[0221] FIG. 17 illustrates a user interface for an asset tree display console function. User interface display 2800 as such as might be caused to display during the processing of block 2340 of FIG. 12, or the processing of method 2500 of FIG. 14, for example; such processing possibly performed by a command, control, and configuration console processor such as 2134 of FIG. 10. User interface display 2800 of FIG. 17 is shown to include system header bar 2402, application information and menu bar 2602, console function header area 2610, asset tree display area 2810, and asset tree node information area 2830. System header bar 2402 and application information and menu bar 2602 are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Console function header area 2610 is as described for the identically numbered element depicted in interface

display 2600 of FIG. 16 with the exception that the current position indicator associated with function/task progress indicator 2614 is shown here advanced to the “Preview” task portion. Asset tree display area 2810 is shown to include area title 2812, “PREVIEW YOUR ASSETS”. In one embodiment, asset tree display area title 2812 displays the name, title, or other identifier for the root node of an asset hierarchy. In one embodiment, asset tree display area title 2812 remains fixed in its position regardless of whether the content displayed beneath is scrolled from its initial or default position. Content displayed in the asset tree display area 2810 beneath area title 2812 is a hierarchical listing of the names, titles, or other identifiers of the nodes in the asset hierarchy beneath the root node. The hierarchical level of a node determines the amount of indentation displayed for its identifier. For example, node entry 2814 for a node named “Sand Filtration System” is a child of the root node, at level 1 in the asset hierarchy, and so displayed with a first amount of indentation. As a second example, node entry 2816 for a node named “Sand Filter 5/6” is a child of the “Sand Filtration System” node, a grandchild of the root node, at level 2 in the asset hierarchy, and so displayed with a second amount of indentation which, in one embodiment, is a greater amount of indentation than for any level above it in the hierarchy. That is to say, the deeper a node is in the asset hierarchy (the higher its level number) the greater the amount of indentation used for its name as displayed in asset tree display area 2810. As yet a further example, node entry 2820 for a node named “Foam Fractionation” is a child of the root node, at level 1 in the asset hierarchy, and so displayed with the same first amount of indentation as for node entry 2814. As one final example, node entry 2822 for a node named “Protein Skimmer 5” is a child of the “Foam Fractionation” node, a grandchild of the root node, at level 2 in the asset hierarchy, and so displayed with the same second amount of indentation as for node entry 2816.

[0222] Node entry 2818 for a node named “Sand Filter 3/4” is shown having a darker background than the other node entries displayed in asset tree display area 2810. The distinctive highlighting of node entry 2818 is a visual indicator that the user has interacted with the interface to indicate a selection of node entry 2818. In an embodiment, a mouse click or touchscreen press in the display area of node entry 2818 may be an indication to the computing machine of the user’s desire to select the particular entry. A selected node entry of asset tree display area 2810, such as node entry 2818, may cause interface 2800 to be refreshed or updated to display information associated with the node represented by the selected node entry in asset tree node information area 2830. In this example, asset tree node information area 2830 is shown to include the area title of “INFORMATION” followed by a list of information items 2832 related to the node which, in an embodiment, may be stored in or in association with the representation of the node in a stored form of the asset tree hierarchy, perhaps as it was just determined. In one embodiment, the list of information items 2832 may include an entry for each column classified as a Metric or Asset Metadata column during the processing associated with interface 2700 of FIG. 16. Each entry in the list of information items 2832 may be displayed as a column or field name followed by the corresponding value. For example information items list entry 2834 is shown with column or field name “Manufacturer.” and the value “Neptune Benson.” In one embodiment, while having an entry in

the list of information items 2832, information items showing a value of “N/A” (indicative of a null value in an earlier search result) are not actually recorded as part of the definitional information in a final stored form of the asset tree hierarchy. Such a final stored form of the asset tree hierarchy may be created and committed to computer storage such as the CCC data store 2132 of FIG. 10, in response to the AMRS receiving an indication of user interaction with “Next” action button 2616 of interface 2800 of FIG. 17. Such processing may be part of the concluding processing contemplated for block 2340 of FIG. 12, in an embodiment.

[0223] FIG. 18 illustrates a user interface display for a metrics console function. User interface display 2900 is such as might be caused to display during the processing of block 2342 of FIG. 12, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. User interface display 2900 of FIG. 18 is shown to include system header bar 2402, application information and menu bar 2902, asset hierarchy display area 2920, metrics overview display area 2940, and metric detail display area 2960. System header bar 2402 is as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Application information and menu bar 2902 is comparable to other application information and menu bars depicted and described in relation to earlier appearing figures, such as bar 2602 of FIG. 17. Asset hierarchy display area 2920 may enable a user to identify one or more assets of an asset tree for which metrics information may be created, edited, deleted, or otherwise processed. Metrics overview display area 2940 may enable a user to view a range of defined metrics associated with an asset hierarchy and to select a particular metric for individual processing, such as editing, or deletion. Metric detail display area 2960 may enable a user to view and interact with information and processing for a particular metric.

[0224] Asset hierarchy display area 2920 is shown to include area title 2922 “Assets”, “Add new asset” action button 2924, asset search component 2926, and asset node list 2928. Asset node list 2928 includes a node list entry for each of one or more of the nodes in an asset hierarchy. In one embodiment, initially and by default, node list 2928 includes a node list entry for every node in the asset hierarchy except for the root node. In an embodiment, a user may interact with asset search component 2926 of interface 2900 to indicate to the AMRS filter criteria as may be applied to the nodes of the asset hierarchy before populating list 2928. For example, a user may enter the word “recovery” into a text box of search component 2926 and the AMRS upon receiving that indication may update asset tree node list 2928 to include entries for only those nodes of the asset tree whose names include the word “recovery.” In such an example, the displayed asset tree node list 2928 would be shortened to show only the “Backwash Recovery” node entry and the node entries for its children, i.e., “Recovery Basin 2” through “Recovery Proteins Skimmer.” The display of asset tree node list 2928 may depict the hierarchical relationship among the nodes using an indentation scheme such as that described in regard to the node list display of 2810 of FIG. 17. Examples of asset node entries of list 2928 shown for interface 2900 include node entry 2930 representing level 1 node “Ozonation System”, node entry 2931 representing level 2 node “Ozone Tower 2”, and node entry 2932 representing level 2 node “Ozone Tower 1”. Asset tree node list

2928 of FIG. 18 may include an interactive check box for one or more node entries to enable a user to identify a selection of one or more asset nodes through interaction with the checkboxes. For example, a user may interact with a number of checkboxes to place them in a selected state, and the set of asset node entries with selected checkboxes may be used to associate the asset tree nodes represented by those entries with a metric being created or edited. In an embodiment, the selection state of checkboxes in the node entries of list 2928 may be set by the AMRS to correspond to the set of assets associated with a metric selected from metrics list 2948.

[0225] A user interaction with “Add new asset” action button 2924 that is indicated to the AMRS may cause the AMRS to engage processing to effect the addition of one or more nodes to the current asset tree hierarchy. In one embodiment, the AMRS may present a user interface similar to interface 2700 of FIG. 16 but including one or more empty table rows that are enabled for data entry by the user such that the user can manually supplement the data of the search result set. Such manual supplementation to the search result set may be useful, for example, in the situation where a user is aware of new assets that are about to come online but for which no machine data has yet been collected. In one embodiment, in response to an indication of user interaction with “Add new asset” action button 2924 of FIG. 18, AMRS may engage processing to effect the addition of one or more nodes to the current asset tree hierarchy by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define an additional node in the asset tree hierarchy. These and other embodiments are possible.

[0226] Metrics overview display area 2940 is shown to include area title 2942 “Metrics”, “Add new Metric” action button 2944, metrics search component 2946, and metrics list 2948. In an embodiment metrics list 2948 may be populated by the AMRS with an entry for each already-defined metric known to the AMRS. In an embodiment, metrics list 2948 may be populated by the AMRS with an entry for each already-defined metric known to the AMRS and associated with at least one node of the asset tree hierarchy represented in whole or in part in 2920. In an embodiment, metrics list 2948 may be populated by the AMRS with an entry for one or more metrics known to the AMRS to be associated with any manufacturer (or other metadata information item) that is associated with any asset node of the asset tree hierarchy represented in whole or in part in 2920. These and other embodiments are possible. In an embodiment, population of metrics list 2948 by the AMRS may be influenced by prior user interaction with metrics search component 2946 to supply a filter criteria.

[0227] Individual entries of metrics list 2948 may each be enabled for user interaction so as to enable a user to indicate a selection of one of the entries of the list, in an embodiment. Metrics list entry 2950 of interface 2900, “HealthScore”, is depicted with a different background color than that of the other entries, indicating a default selection of the metric associated with that entry or indicating a prior user interaction with that entry to effect such a selection. The selection of an entry of metrics list 2948 may result in the display of information and user interface elements in a metric information area 2960 that are pertinent to the metric represented by the selected entry, such as selected entry 2950. User interaction with “Add new Metric” action button 2944 may

be indicated to the AMRS, which may in response present a modified or alternate user interface display enabling a user to indicate sufficient information to define a new metric.

[0228] Metric detail display area 2960 displays information and user interface elements pertinent to the metric represented by the selected entry of metrics list 2948, such as selected entry 2950, “HealthScore”. Metric detail display area 2960 is shown to include general information area 2962, metric configuration action area 2964, alerts configuration action area 2970, actions configuration action area 2974, and metric deletion action area 2980. General information area 2962 may display one or more information items of the defined metric using a fieldname: value format (e.g. “Metric Name: Healthscore”), and may display one field or information item on each line. Metric configuration action area 2964 is shown to indicate a count of the number of assets (asset nodes) associated with the subject metric, which count may correspond to the number of node entries of list 2928 having selected checkboxes, in an embodiment. Metric configuration action area 2964 is shown to include Edit action button 2966 enabling a user to indicate a desire to view and change information of a metric configuration. Alerts configuration action area 2970 is shown to include an area title or caption, and Edit action button 2972. Actions configuration action area 2974 is shown to include an area title or caption, and Edit action button 2976. In response to an indication of user interaction with any of Edit action buttons 2966, 2972 and 2976, the AMRS may cause the display of an updated, altered or alternate user interface that displays appropriate information and enables a user to indicate changes thereto. Examples of such are illustrated and discussed in relation to figures that follow.

[0229] Metric deletion action area 2980 is shown to include an area title or caption, a deletion warning message, and “Delete Metric” action button 2982. In response to an indication of user interaction with action button 2982, the AMRS may display a user interface element, such as a message box, requiring a final user confirmation for a deletion action and, upon receipt of an indication of such a confirmation, the AMRS may engage processing to erase, suspend, deactivate, flag, mark, or otherwise logically or physically “delete” the metric definition. Such processing may further include updating metrics list 2948 to remove the entry representing the deleted metric, in an embodiment.

[0230] FIG. 19 illustrates a user interface display for a metrics configuration console function. In an embodiment of an AMRS, interface 3000 of FIG. 19 may be caused to be displayed on a user interface device in response to user interaction with “Add new Metric” action button 2944 or Edit action button 2966 of FIG. 18. When interface 3000 of FIG. 19 is displayed in response to user interaction with “Add new Metric” action button 2944 of FIG. 18, information components of the user interface may be empty, contain default values, and/or display tips to the user regarding the use of the component. When interface 3000 of FIG. 19 is displayed in response to user interaction with Edit action button 2966 of FIG. 18, information components of the user interface may be populated with information values from the existing metric definition. Interface 3000 of FIG. 19 is shown to include header area 3002, footer area 3004, general information area 3010, data selection area 3020, metric determination area 3030, and metric time factor area 3060. Header area 3002 of interface 3000 is shown to include a title or caption for the interface.

[0231] General information for the metric definition is displayed and/or indicated by the user via the interface components of general information area **3010**. General information area **3010** is shown to include a metric name component **3012**, a metric description component **3014**, and a metric units components **3016**. In this illustrative example, the metric name component **3012** is an interactive text box used by the AMRS to display information it has for the name or identifier of the metric and to receive indications of a user-desired value for the name or identifier of the metric. Similarly, the metric description component **3014** is an interactive text box used by the AMRS to display information it has regarding a description for the metric and to receive indications of a user-desired value for the description. Similarly, the metric units component **3016** is an interactive text box used by the AMRS to display information it has for a designated unit to be associated with values identified or produced for the metric and to receive indications of a user-desired value for the metrics unit. For example, a metric unit may be a unit of measurement such as degrees-Fahrenheit or pounds-per-square-inch (PSI).

[0232] In one embodiment, metrics may be determined by search queries executed against asset-related data (that may include other metric data). In one embodiment, the asset-related data are kept by and/or accessed via a data intake and query (DIQ) system such as DIQ **108** of FIG. 1 or DIQ **2120** of FIG. 10. In one embodiment, more particularly, asset-related metrics are determined by search queries processed by a DIQ against data it stores, indexes, manages, and/or provides access to for search processing. Accordingly, in such an embodiment, metric configuration or definition data may relate directly to the specification of search queries to be executed using the processing facilities of a DIQ. Such search queries may be executed by a DIQ just as the definition-time search discussed earlier in relation to block **2514** of FIG. 14, but are, in contrast, metric- or metric-time search queries for deriving metric data/measurements. In that vein, data selection area **3020** of FIG. 19 is used to display or specify information of a metric definition related to data selection aspects of a search query for the metric. Data selection area **3020** of interface **3000** is shown to include a set of option buttons **3022** including an “Asset Model” option button **3023a** and an “Ad hoc search” option button **3023b**. Data selection **3020** is shown to further include chosen-asset component **3026** and aggregation mode selection radio buttons **3027a-b**. In an embodiment, when “Asset Model” option button **3023a** is activated (to the exclusion of option button **3023b**) a user is enabled to indicate to the AMRS a selection of one or more assets represented by nodes in the asset hierarchy, and the AMRS processes the indicated selection to determine a search query, aspects, or portions that cause the executed search query to match, select, filter, or identify the data properly associated with the asset selection and as may be relevant to the production of the metric value.

[0233] In one embodiment, a user is enabled to indicate a selection of assets by entering their identifiers into chosen-asset component **3026**. In one embodiment, a user is enabled to indicate a selection of assets by selecting one or more assets from a displayed list of nodes in the asset hierarchy. In such an embodiment, the AMRS when processing the user indications may populate chosen-asset component **3026** with the names of the assets that were indicated for selection by the user.

[0234] Aggregation mode selection radio buttons **3027a-b** may not strictly be related to a data selection aspect of a DIQ search query but may be related to an aspect of a DIQ search query that determines how the selected data is processed to produce a search query result. In one embodiment, indicating individual (non-aggregated) metric mode to the AMRS using radio button **3027a** may result in the AMRS determining a search query, aspect, or portion that causes the executed search query to produce a value for the metric on a per-asset basis, instead of or in addition to producing the value for the metric on an aggregated basis. In one embodiment, indicating aggregated metric mode to the AMRS using radio button **3027b** may result in the AMRS determining a search query, aspect, or portion that causes the executed search query to produce a value for the metric on an aggregated basis. In an embodiment, a search query that does not indicate a need for per-asset results (or that indicates only an aggregated result is required) may permit the DIQ to execute the search query in an optimized fashion with early data aggregation, reducing resource consumption and speeding the execution of the query. Accordingly, an embodiment enabling the specification of a more efficient option when it will accommodate the needed result is an improved data processing machine.

[0235] Metric determination area **3030** may be used to display or specify information of a metric definition related to aspects of a search query that determine a value for the metric from the data selected by the search query. Metric determination area **3030** of interface **3000** is shown to include calculation category component **3032**, calculation line deletion action element **3034**, metric component **3036**, calculation component **3050**, filter component **3052**, units component **3054**, operator component **3056**, “Add new logic” action element **3058**, and “Save logic” action element **3048**.

[0236] Calculation category component **3032** is illustrated in this example as a drop-down selection box indicating a default or last-user-selected value of “New metric logic” as the calculation category. Other calculation categories made available by drop-down selection box **3032** may include, for example, “Built-in Metrics” and “Advanced Analytics”. In one embodiment, the determination of a metric value in the “New metric logic” calculation category may be specified in terms of one or more calculation lines. Components **3034**, **3036**, **3050**, **3052**, **3054**, **3056** of metric determination area **3030** together represent such a calculation line. Additional calculation lines may be displayed by the AMRS as the result of processing user indications of a desire for such additional calculation lines by interaction with “Add new logic” action element **3058**, in one embodiment. In one embodiment, calculation line deletion action element **3034** enables a user to indicate a desire to delete the calculation line of which is part. The AMRS receiving such an indication may perform such a deletion and refresh or update the interface display **3000**, accordingly.

[0237] Metric component **3036** is a drop-down selection box having an indication of the current selection **3038**, a selection deletion/cancellation action component **3042**, and a drop-down display action button **3040**. In one embodiment, the selection list of drop-down selection box **3036** may be populated with the names of all metric fields/columns associated with any one or more nodes of the asset hierarchy associated with the metric being defined. In one embodiment, the selection list of drop-down selection box

3036 may be populated with the names of all metric fields/columns that is each associated with all of the nodes of the asset hierarchy associated with the metric being defined. Other embodiments are possible. Calculation component **3050** is shown as a drop-down selection box displaying an indication of the currently selected calculation, here shown as “Average”. The calculation option indicated by calculation component **3050** is the calculation or other processing performed over the values for the field identified at **3036** in the selected data during the execution of the search query. An entry of filter criteria at **3052** may be used to limit the calculation of **3050** over the values of the field of **3036** for something less than all of the data selected using the criteria of **3020** alone.

[0238] An entry of a units value and units component **3054** may be utilized by the AMRS to tag or label calculation results internally or for external displays, or may be utilized by the AMRS for normalization of different values to a common unit. Operator component **3056** is a drop-down selection list enabling the user to make a selection of an operation to be performed between the calculation of the calculation line in which operator component **3056** appears and a subsequent calculation line as may be added by interaction with action element **3058**. In one embodiment, the selection list of operator component **3056** may be populated with operators including AND, OR, +, -, /, x, and %, for example.

[0239] Metric time factor area **3060** may be used to display or specify information of a metric definition for time-related aspects of the search query. Metric time factor area **3060** is shown to include schedule component **3062** and span component **3064**. Span component **3064** is shown as a drop-down selection list displaying “per hour” as the current selection by default, prior user interaction, or other means. The AMRS processes an indicated selection of span component **3064** to determine a search query, aspects, or portions that cause the executed search query to match, select, filter, or identify the data pertaining only to the selected time span. While the indicated value for the span component **3064** is, in a sense, internal to the search query, affecting the processing an execution of the search query will perform, in one embodiment, an indicated value for schedule component **3062** is, in a sense, external to the search query, not affecting the processing an execution of the search query will perform, but rather affecting when and/or how often the search query is performed. Schedule component **3062** is shown as a drop-down selection list displaying “1 minute” as the current selection by default, prior user interaction, or other means. In one embodiment, the drop-down selection list of schedule component **3062** may include 1 minute, 5 minutes, 10 minutes, 15 minutes, 20 minutes, 30 minutes, 1 hour, 2 hours, 4 hours, 6 hours, 12 hours, 1 day, 2 days, 3 days, and 7 days, as available options, for example.

[0240] Footer area **3004** of interface **3000** is shown to include Cancel action button **3070**, Next action button **3072**, and Finish action button **3074**. In response to receiving an indication of user interaction with Finish action button **3074**, the AMRS of one embodiment may create and store a proper representation of metric definition in computer storage, including representations of associations between the metric and one or more nodes of the asset tree, and may cause navigation to another user interface, perhaps interface **2900** of FIG. **18**. In response to receiving an indication of user interaction with Next action button **3072** of FIG. **19**, the

AMRS of one embodiment may create and store a proper representation of metric definition in computer storage, including representations of associations between the metric and one or more nodes of the asset tree, and then cause navigation to a user interface for configuring conditions and alerts based on the metric, such as the interface illustrated and described in relation to FIG. **20**.

[0241] FIG. **20** illustrates a user interface display for a metrics condition and alerts console function. In an embodiment of an AMRS, interface **3100** of FIG. **20** may be caused to be displayed on a user interface device in response to user interaction with “Next” action button **3072** of FIG. **19** or Edit action button **2972** of FIG. **18**. When interface **3100** of FIG. **20** is displayed in response to user interaction with “Next” action button **3072** of FIG. **19** the data information components of the user interface may be empty, contain default values, and/or display tips to the user regarding the use of the component. When interface **3000** of FIG. **19** is displayed in response to user interaction with Edit action button **2972** of FIG. **18** the data information components of the user interface may be populated with information values from an existing condition/alert definition. Interface **3100** of FIG. **20** is shown to include header area **3102**, footer area **3104**, condition name component **3110**, condition description component **3112**, condition line deletion action component **3116**, condition operator component **3120**, threshold/multiplier component **3122**, urgency component **3124**, alert option component **3126**, “Add new condition” action element **3118**, “Save conditions” action element **3128**, condition schedule component **3132**, and condition span component **3134**. Header area **3102** of interface **3100** is shown to include a title or caption for the interface.

[0242] In this illustrative example, the condition name component **3110** is an interactive text box used by the AMRS to display information it has for the name or identifier of the condition and to receive indications of a user-desired value for the name or identifier of the condition. Similarly, the condition description component **3112** is an interactive text box used by the AMRS to display information it has regarding a description for the condition and to receive indications of a user-desired value for a description for the condition.

[0243] Footer area **3104** of interface **3100** is shown to include Cancel action button **3142**, Back action button **3144**, Next action button **3146**, and Finish action button **3148**. In response to receiving an indication of user interaction with Finish action button **3148**, the AMRS of one embodiment may create and store a proper representation of a condition definition in computer storage, and may cause navigation to another user interface, perhaps interface **2900** of FIG. **18**. In response to receiving an indication of user interaction with Back action button **3144** of FIG. **20**, the AMRS of one embodiment may create and store a proper representation of a condition definition in computer storage and then cause navigation to a prior user interface. In response to receiving an indication of user interaction with Next action button **3146**, the AMRS of one embodiment may create and store a proper representation of a condition definition in computer storage and then cause navigation to a user interface for configuring actions based on the condition, such as the interface illustrated and described in relation to FIG. **21**.

[0244] In one embodiment, the determination of a condition may be specified in terms of one or more condition lines, much as the determination of a metric value may be

specified in terms of calculation lines as described in reference to FIG. 19. Components **3116**, **3120**, **3122**, **3124**, and **3126** of FIG. 20 may together represent such a condition line. Additional condition lines may be displayed by the AMRS as the result of processing user indications of a desire for such additional condition lines by interaction with “Add new condition” action element **3118**, in one embodiment. In one embodiment, condition line deletion action element **3116** enables a user to indicate a desire to delete the condition line of which is part. The AMRS receiving such an indication may perform such a deletion and refresh or update the interface display **3100**, accordingly.

[0245] Condition operator component **3120** is a drop-down selection box having an indication of the current selection, here shown as “Greater than (>),” which may be a default or last-user-selected value. Inasmuch as the value selected for the condition operator component is a comparison operator, its first comparand is the value of the metric with which it is associated, and its second comparand is the value displayed and/or specified via threshold/multiplier component **3122**. Threshold/multiplier component **3122** is depicted in interface **3100** as an editable text box enabling a user to indicate new values or edit those previously displayed by the AMRS. During operation of the AMRS to perform asset monitoring and reporting in accordance with a condition and alert definition established by the use of interface **3100**, a determination may be made by the AMRS that a condition specified by the combination of a metric value, a condition operator, and a threshold/multiplier value, has been met. As a result of that determination, the AMRS may perform or cause actions that may include recording a condition event record, as one example. In an embodiment, actions performed or caused by the AMRS as the result of such a determination may be characterized, classified, or attributed with an urgency level displayed and/or indicated at definition time by urgency component **3124** of interface **3100**. In the example where a condition event record is recorded upon determination of the condition, the event record may include an urgency field having a value taken from the condition definition. In one embodiment, where multiple conditions are defined for a metric using multiple condition lines of an interface such as **3100**, the conditions may be effectively evaluated for the metric in the order in which they were defined, and a first successful condition determination may terminate the search for any other satisfied condition. In one embodiment, where multiple conditions are defined for a metric using multiple condition lines of an interface such as **3100**, the conditions may each be evaluated in a particular instance without regard to the success of any other, and multiple conditions may be determined to be satisfied and may be processed as such. These and other embodiments are possible.

[0246] Alert option component **3126** is shown for interface **3100** as a binary set of mutually exclusive selection buttons. Other embodiments are possible. The selection indicated by alert option component **3126**, whether by default or by user interaction, is reflected in the definition information for the condition.

[0247] During operation of the AMRS to perform asset monitoring and reporting in accordance with a condition and alert definition established by the use of interface **3100**, a determination may be made by the AMRS that a condition specified by the combination of a metric value, a condition operator, and a threshold/multiplier value, has been met. As

a result of that determination, the AMRS may perform one or more alert actions which performance, itself, may be conditioned directly or indirectly on the alert option value reflected in the definition information for the condition. In one embodiment, an alert action may include posting an alert event to an event data store. In one embodiment, an alert action may be an action defined through the use of a user interface such as that depicted and described in relation to FIG. 21.

[0248] In one embodiment, an AMRS may utilize the processing of a DIQ to perform search queries against metrics data to make determinations regarding the satisfaction of one or more conditions. In such an embodiment, interface components **3132** and **3134** of interface **3100** may be used to display and indicate time related values for such a search query. Condition schedule component **3132** may be used to display and indicate a desired value for a frequency with which to schedule the condition satisfaction search query. Condition span component **3134** may be used to display and indicate a desired value that determines the duration of the time span of data considered by an execution of the condition satisfaction search query.

[0249] FIG. 21 illustrates a user interface display for an actions console function. In an embodiment of an AMRS, interface **3180** of FIG. 21 may be caused to be displayed on a user interface device in response to user interaction with “Next” action button **3146** of FIG. 20 or Edit action button **2976** of FIG. 18. Interface **3180** of FIG. 21 is shown to include header area **3182**, footer area **3184**, and for action-type elements **3192**, **3194**, **3196**, and **3198**. Header area **3182** is shown to display a title or caption for the interface. Each of action-type elements **3192**, **3194**, **3196**, and **3198** represents a defined action of a particular type that may be performed, initiated or caused by processing of the AMRS in association with condition/alert processing such as may be defined or configured via a user interface such as interface **3100** FIG. 20. The defined actions may be parameterized or otherwise configurable or customizable. In one embodiment, a defined action may be packaged as a form of plug-in or extension module for the AMRS, and such a module may include sufficient computer instructions, code, scripts, or the like, media properties (e.g., icons), and any other related data, to enable the use of the defined action in the AMRS environment. In one embodiment, user interaction with an action-type element may result in the AMRS simply activating the defined action and possibly updating interface display **3180** to indicate the activation, perhaps by changing some aspect of the visual appearance of the corresponding action-type element (e.g., background color). Such processing may be useful in the case of the defined action that has no customizable or configurable parameters, or that has customizable or configurable parameters the values for which can be automatically determined using system or user profile information, for example. In one embodiment, user interaction with an action-type element may result in the AMRS causing the display of a different user interface or additional user interface components that enable a user to indicate customization or configuration information for the defined action instance. These and other embodiments are possible. Example action-type elements shown for interface **3180** include action-type element **3192** corresponding to a defined action of sending an SMS message (specifically via the Twilio service), action-type element **3194** corresponding to a defined action of sending an e-mail message (perhaps

using IMAP), action-type element **3196** corresponding to a defined action of creating a work order in an incident tracking system (here, specifically a ServiceNow® system), and action-type element **3198** corresponding to a defined action of a change in lighting color (here, specifically via Philips Hue Lights control technology).

[0250] Footer area **3184** is shown to include Cancel action button **3186**, Back action button **3187**, and Finish action button **3188**. In response to receiving an indication of user interaction with Finish action button **3188**, the AMRS of one embodiment may create and store a proper representation of an action definition in computer storage, and may cause navigation to another user interface, perhaps interface **2900** of FIG. 18. In response to receiving an indication of user interaction with Back action button **3187** of FIG. 21, the AMRS of one embodiment may create and store a proper representation of an action definition in computer storage and then cause navigation to a prior user interface.

[0251] FIG. 22 illustrates a user interface display for creating or editing a custom monitoring or reporting presentation for an asset tree. Interface **3200** of FIG. 22 is such as might be caused to display during the processing associated with block **2346** of FIG. 12, for example. In one embodiment, a custom monitoring or reporting presentation for asset tree information may be termed a display or presentation view, and a view instance produced by the AMRS during the performance of its monitoring and reporting processes may be based on a view template that is configured and/or customized by the user. In one embodiment, such a view may be considered a dashboard, and its template a dashboard template. Interface **3200** of FIG. 22 illustrates one possible embodiment of a user interface whereby a user may create, configure, and/or customize such a view template.

[0252] Interface **3200** of FIG. 22 is shown to include system header bar **2402**, application information and menu bar **2902**, function header bar **3206**, asset display area **3230**, toolbar **3222**, view template display area **3250**, and configuration display area **3270**. System header bar **2402** and application information and menu bar **2902** are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures.

[0253] Function header bar **3206** is an embodiment of a header bar as may be useful in an interface associated with the function of creating, configuring, and/or customizing a view template. Function header bar **3206** is shown to include title **3210**, display mode action button **3212**, timeframe component **3214**, Clear action button **3216**, Revert action button **3218**, and Save action button **3220**. Except for title **3210**, interface components of function header bar **3206** of this illustrative example are interactive elements that enable a user to make indications of data values and desired actions, for example, to AMRS functionality, which the AMRS computing machinery can then process according to its design, circuitry, and programming.

[0254] Display mode action button **3212** enables the user to request a toggling action between two alternate display modes. An editing display mode is represented in interface **3200** of FIG. 22 as it appears and provides user interface components for creating, configuring, and/or customizing (i.e., editing) a view template. When in editing display mode, display mode action button **3212** displays the name of the alternate display mode, i.e., "View". A user interaction

with display mode action button **3212** may result in the transition to a View mode user interface display where a full-screen or near full-screen view is presented based on the current working state of the view template under construction. Timeframe component **3214** may be a drop down selection list interface component that enables a user to indicate a desired time or time frame of data to use in relation to data-driven or data-aware elements that may be included in a view template. A user interaction with Clear action button **3216** in an embodiment may cause the view template under construction to be emptied of all of its content. A user interaction with Revert action button **3218** in an embodiment may have the effect of causing recent changes made to a view template to be abandoned. In one embodiment a user interaction with Revert action button **3218** may cause the abandonment of only the single most recent change to the view template. In one embodiment, a user interaction with Revert action button **3218** may cause the abandonment of all changes made to the view template since the last time Save action button **3220** was activated. In one embodiment, a user interaction with Revert action button **3218** may cause the abandonment of all changes made to the view template since the last time an autosave action was performed by the AMRS console processor. Other embodiments are possible. A user interaction with Save action button **3220**, in one embodiment, may cause the current configuration of the view template under construction to be reflected in computer storage anyway such that it may be recalled or restored, perhaps by reflecting the configuration information in a named file in the filesystem of a host computer.

[0255] Asset display area **3230** is shown to include area title **3232**, asset search component **3234**, and asset node list **3236**. Asset node list **3236** has similar content, organization, appearance, and formatting to asset node lists of user interfaces illustrated and discussed in relation to earlier appearing figures. Similarly, a working understanding of asset search component **3234** may be developed by consideration of what has come before. Asset node list **3236** is shown to include asset node list entry **3240** representing the level 1 asset node named "Output", asset node list entry **3242** representing a child of the "Output" node, named "Output Pump-3", and asset node list entry **3244** representing a child of the "Output Pump-3" node, named "Flow". "Flow" asset list node entry **3244** is shown as having a darker background color than the other asset node list entries, indicating that it is the currently selected asset node list entry.

[0256] Toolbar **3222** is shown as having a number of tool icons such as "T" icon **3224**. Icons in toolbar **3222** may be selected by a user interaction to cause a particular effect, engage a particular function, and/or begin a particular operational mode. For example, in one embodiment, a mouse click on icon **3224** may cause the addition of an empty, default-sized text display element to the view template under construction, causing it to appear at a default location in view template display area **3250**, and engaging an operational mode for entering text into the newly introduced text display element. In an embodiment, many of the tools represented in toolbar **3222** may be associated with adding different types of elements to the view template, and with manipulating the elements that are present in the template. In one embodiment, one or more tool icons may be associated with static elements that may be included in the design of the view template. Such static elements may include text

blocks or labels, imported graphical imagery (e.g., icons, picture files, videos, fixed animations), or drawing elements such as shapes and lines. In one embodiment, one or more tool icons may be associated with dynamic elements that may be included in the design of the view template. Such dynamic elements may be data-driven or data-aware and may determine one or more aspects of their appearance or behavior at a point in time based on currently supplied data. Such data-driven or data-aware dynamic elements may be referred to as “widgets” in one embodiment. In one AMRS embodiment, a monitoring/reporting processor, and CCC console processor functions related thereto, may include functionality to implement a number of built-in widgets and may further include functionality to implement an extensible widget framework which functionality may include, for example, functionality to recognize, install, or activate widget modules, and functionality to exercise the content of those modules. In one such embodiment, widget modules may be packaged after the fashion of programming objects and have attributes or properties (associated data) and methods or behaviors (programmed actions) which may be accessible and/or exercisable by a recognized interface. In one embodiment, a supported widget may be limited to receiving a single data factor or component that drives it, such as the data of a particular metric for a particular asset. In one embodiment, a supported widget may be able to receive multiple data factors or components to drive it, such as the data of different metrics that may be associated with the same asset. These and other embodiments are possible.

[0257] View template display area 3250 is shown to include a representation of some or all of the view template currently under construction. (For example, a representation of only some of the current view template may appear in display area 3250 where display area 3250 is smaller than the size of the current view template. In such a case, view template display area 3250 may be scrollable.) The current view template is shown to include static graphical elements including, for example, pipe segment icon/shape/picture 3252. The current view template is further shown to include multiple dynamic elements including widgets 3254, 3256, and 3260. Widget 3254 is used to depict an Ozone-Tower asset and presents a graphical depiction of a tank or tower. Data of a fill-level metric associated with the Ozone-Tower asset may drive the appearance of the widget, for example, by adjusting the position of the boundary 3254a between a lower, blue, fill portion and an upper, black, empty portion of the tank/tower depiction. Widget 3256 is used to depict a protein skimmer (PSK) asset and presents a graphical depiction of a protein skimmer apparatus with three metric presentation blocks or tiles beneath it. Data related to a Water-level metric, an ORP-Level metric, and a Temperature metric may drive the appearance of the widget, and, particularly, the current value for each of the metrics is displayed in a corresponding one of the metric presentation blocks or tiles of the widget, and an urgency level associated with each of the values determines the color of the text used to display the current value in the metric presentation blocker tile (for example, the value of 100 for the Water-level metric may be associated with a critical urgency level and so may display in red). In one embodiment, the metric presentation blocks or tiles discussed in relation to widget 3256 may each be an independent widget.

[0258] Widget 3260 is used to depict an asset named “Output-Pump-3”. The appearance of the bounding box with

control points around widget 3260 may indicate that widget 3260 is the active element of the current view template. In one embodiment, the active element of the current view template may have its detailed configuration information presented in configuration display area 3270. Widget 3260 is shown to include a title 3262 that corresponds to the name 3274a of the asset designated as the assigned asset of the widget as seen in the assigned asset area 3274 of configuration display area 3270. Widget 3260 is further shown to include pump icon 3264 which comes from a Pump.svg file as seen in icon area 3272 of configuration display area 3270. Widget 3260 is further shown to include metric display block/tile 3266 which is driven by the data for a Flow metric as seen by metric token 3276a in the Display Metrics area 3276 of configuration display area 3270. Configuration display area 3270 is shown to further include Display Alerts area 3278 where defined alerts to be displayed in association with the widget are represented by tokens, such as 3278a; general drawing attributes area 3279; title display option area 3280; data-driven animation control section 3282; drilldown area 3284 where a user interface navigation target 3284a, such as a Diagnose interface display, is defined to use in circumstances where a user double clicks or performs some other specified interaction with the displayed widget; and update action button 3286 which enables a user to indicate the desire to synchronize the representation of the widget displayed in 3250 with the representation of its configuration information displayed and 3270.

[0259] It is noted that a widget-depicted asset may have a corresponding asset node list entry visible in asset display area 3230, or not—as in the case where a superior node of the asset is in a collapsed state in asset node list 3236. For example the Ozone-Tower asset represented by widget 3254 may be subordinate in the asset hierarchy to the high level Ozonation-System asset represented by list entry 3246, but because the Ozonation-System list entry 3246 is in a collapsed state, as indicated by the “>” character in the entry, a list entry for the Ozone-Tower is suppressed and not visible.

[0260] When the creation, construction, editing, or such for a view template is complete, a user may activate Save action button 3220 to safely fix a definition/configuration of the view template in computer storage and exit interface 3200.

[0261] The user interfaces already discussed in relation to asset hierarchy monitoring and reporting, namely, the user interfaces illustrated and discussed in relation to FIGS. 13, 15, 16, 17, 18, 19, 20, 21, and 22, have largely related to interfaces employed by a command, control, and configuration console processor (such as CCC console 2134 of FIG. 10). These interfaces enable a user, such as a system administrator or operator, to manipulate the virtual levers, buttons, dials, and switches (embodied in the information of a CCC data store such as 2132 of FIG. 10) that control the operation of the asset monitoring and reporting machine. The focus now turns with the discussion of the figures that follows to the output side of the work performed by the AMRS machinery. Figures that follow relate largely to reporting interfaces that may be utilized by the AMRS during the processing described and contemplated for block 2350 of FIG. 12 as performed by a monitor/reporter processor such as 2142 of FIG. 10, for example.

3.3 Asset Hierarchy Monitoring/Reporting

[0262] FIG. 23 illustrates a user interface display of a custom asset tree presentation. Interface 3400 of FIG. 23 may be used with great effect to report against a defined asset tree hierarchy with a dynamic and highly customized presentation based on a view template as may have been created using an interface such as 3200 of FIG. 22. Interface 3400 FIG. 23 is shown to include system header bar 2402, application information and menu bar 2902, tab controls area 3410, and active tab display area 3420. System header bar 2402 and application information and menu bar 2902 are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Tab controls area 3410 is shown to include three tab controls, 3412, 3414, and 3416, each corresponding to a different presentation model for asset related data. Tab interfaces are understood in the art and will not be elaborated here. Tab 3412, "Spatial View" is shown to be the active tab.

[0263] Active tab display area 3420 is shown to include spatial view header area 3422, and main display area 3424. Spatial view header area 3422 is shown to include a title indicative of the presentation model associated with the tab, but in another embodiment may additionally or alternatively include a title indicative of the asset tree being reported or indicative of the particular view template used to generate the display content. Other embodiments are possible. Spatial view header area 3422 is further shown to include Edit action button 3428 which may enable a user to navigate to an interface such as 3200 of FIG. 22 in order to modify the view template on which the presentation of user interface 3400 of FIG. 23 is based. Main display area 3424 presents a display based on a view template and instantiated using current data. Monitor/reporter processor 2142 of FIG. 10 may, on some continuous basis, update or refresh the content of main display area 3424 using updated or refreshed data.

[0264] FIG. 24 illustrates a user interface display for a metrics view. Interface 3500 of FIG. 24 is shown to include system header bar 2402, application information and menu bar 2902, tab controls area 3410, metrics view header area 3510, and main display area 3520. System header bar 2402, application information and menu bar 2902 and tab controls area 3410 are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Tab 3414, "Metrics View" is shown to be the active tab. Main display area 3520 is shown to make a presentation of assets related metrics data in a tabular format. Main display area 3520 is shown to include column header 3522 and table data area 3524. Column header 3522 is shown to include column headings 3522a-e which correspond respectively to the columns of the table 3540a-e. Column heading 3522a displays "Asset" as the heading for first column 3540a. Data in column 3540a for each of table rows 3550a-s is the name or identifier of an asset node to which all of the data appearing in the row pertains. The remainder of the column headings 3522b-e each contain the name of the metric for which measurement data appears in the column, and an indication of the relevant "Unit" or unit-of-measure designation that denominates the measurement values for the metric. Column header 3522 is further shown to include an add-column action button 3522z. Interaction with ad-column action button 3522z may cause the appearance of another data column in the tabular display and a selection list populated

as with the names of available metrics associated with nodes of the asset hierarchy. A user may indicate a selection from the list and the monitor/reporter processor of the AMRS will place the name of the indicated selection, along with the appropriate Unit value in the column heading of the new column and will populate measurement for that data into each row of the table, as available. When populating measurement data for the metrics into the rows of the table, generally, an embodiment may further determine a visual attribute or ornament for each cell that corresponds to an urgency level associated with the value of the measurement for the metric represented in the column. In the illustrative embodiment of FIG. 24, the background color of a table cell is the visual attribute used for indicating the urgency level associated with the measurement value. Accordingly, different background colors can be seen in the cells and, for example, the 5% measurement value for a Downtime metric as seen in cell 3560 may be considered a normal urgency value and so display with the background color of green, while the 7% measurement value for the Downtime metric as seen in cell 3562 may be considered a warning urgency value and so display with a background color of yellow, while the 9% measurement value for a Downtime metric as seen in cell 3564 may be considered a critical urgency value and so display with the background color of red. Notably, cell 3566 displays with a background color of gray indicating that there is no measurement data for the reported time period, for the metric indicated by the column of cell (3540e) and the asset indicated by the row of the cell (3550q). Such a condition may exist, for example, where the column metric is not relevant to the row asset, or where machine data used for the measurement was not produced, lost, or delayed.

[0265] FIG. 25 illustrates a user interface display for a conditions and alerts view. Interface 3600 of FIG. 25 is shown to include system header bar 2402, application information and menu bar 2902, tab controls area 3410, conditions and alerts view header area 3610, and main display area 3620. System header bar 2402, application information and menu bar 2902, and tab controls area 3410 are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Tab 3416, "Conditions & Alerts View" is the active tab. Main display area 3620 is shown to make a presentation of conditions and alerts data in a tabular format. Main display area 3620 is shown to include column header 3622 and table data area 3624. Column header 3622 is shown to include column headings respectively corresponding to the columns of the table 3640a-h, namely, "Condition Name", "Category", "Timestamp", "Source", "Asset", "Metric", "Value", and "Unit". The rows of table data area 3650a-p are by default populated with conditions and alerts event information in reverse chronological order, i.e., in descending order under Timestamp column heading 3622c. The background color of each row is determined by the urgency category associated with the row. Rows with "Normal" appearing in their "[urgency] Category" column have a green background; rows with "Medium" appearing in their "[urgency] Category" column have a yellow background; and rows with "Critical" appearing in their "[urgency] Category" column have a red background.

[0266] FIG. 26 illustrates a user interface display for a diagnostics view. Interface 3700 of FIG. 26 may provide a deep-dive or drill-down interface which may be a destina-

tion of choice for an analyst who observes or suspects a problem based on information in a higher-level and/or less detailed display, and then needs to see the detail in order to diagnose the problem. As one example, a user may regularly monitor her water filtration system using the spatial view presentation of interface 3400 of FIG. 23. If at some point, its urgency category becomes critical, and the widget representing the Sand Filter pump goes red, the spatial view interface may be so configured that the user may interact with the Sand Filter pump, such as by double-clicking, to navigate directly to interface 3700 of FIG. 26. Interface 3700 is shown to include system header bar 2402, application information and menu bar 2602, diagnose header 3710, graph lanes 3720, 3630, 3640, and 3650, and viewport time range indicator 3762. System header bar 2402 and application information and menu bar 2602 are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Diagnose header 3710 is shown to include asset identifier 3712, new lane selector 3714, timeframe selector 3716, comparison timeframe selector 3718, “Send alert” action button 3764, and “Save” action button 3766. In one embodiment, user interaction with new lane selector 3714 may cause the display of a selection list from which the user may indicate a selection of an available metric or other data source to populate an additional graph lane in interface 3700. In one embodiment, user interaction with timeframe selector 3716 may cause the display of a list of options for timeframe from which the user may indicate a desired selection. In response to receiving the user’s indication of a selection the AMRS may cause search query executions to gather the data for the newly selected timeframe for each of the graph lanes displayed in the interface the AMRS may then update or refresh interface 3700 with the new data and may update viewport time frame indicator 3762 appropriately. In one embodiment, user interaction with comparison timeframe selector 3718 may cause the display of a list of options from which the user may select a timeframe of data to plot in the graph lanes for comparison against the primary timeframe as designated at 3716. Options in the selection list associated with 3718 may be absolute or relative timeframe options and embodiments may vary as to the manner in which the comparison data is presented in the graph lanes, for example, whether side-by-side, over-under, or superimposed. A user may interact with “Save” action button 3766 to save the current configuration and settings of interface 3700 as a template or pattern to reuse at some future date.

[0267] In one embodiment, each of the graph lanes of the diagnostics interface has a graph lane information and control portion and a graph viewport portion. For example, graph lane 3720 has graph lane information and control portion 3722 and graph viewport portion 3724. A graph lane information and control portion may include the name of the metric or other data source that provides the data represented in the graph viewport portion. Measurement data of the metric or data from another source may then be plotted or otherwise graphed or represented in the graph viewport portion of the graph lanes. An axis of the graph viewport portion represents a time dimension (often the long axis) and the graph lane data is depicted as a time series. In one embodiment, a common time axis is used for all of the displayed graph lanes. In one embodiment, the non-time axis may be auto scaled by AMRS processing determinations. In one embodiment, a Settings icon such as 3723 may be

available for each of the graph lanes and enable a user to adjust various settings for the graph lane, for example, the graph style or color. These and other embodiments are possible.

[0268] FIG. 27 illustrates a user interface display for a map view of asset tree data. Interface 3800 of FIG. 27 may make the map view visualization available to the user via a tab control 3802. The map view of 3800 is shown to include map header area 3810 displaying an name, title, or caption for the map, perhaps a name associated with the root node of the asset hierarchy. Map header area 3810 is further shown to include “Edit” action button 3814. User interaction with Edit action button 3814 in one embodiment may result in the presentation of a different or updated interface that enable the user to modify the content, layout, configuration, and definition of the map view template. A map view template may include static elements such as the image of the background map and may include dynamic, data-driven elements such as asset icons 3830, 3832, and 3834, for example, and such as metric widgets or tiles 3822a-c appearing in a metrics reporting area 3822. Defined metrics and conditions may be used to determine an urgency category and or one or more particular attributes of an asset icon in the map view, such as its color. User interaction with an asset icon such as a hover-over interaction with asset icon 3830 in one embodiment may result in the transient display of an asset information box 3840. Embodiments may vary greatly as to the amount and types of information presented in an asset information box which may include, for example, asset definition metadata, associated metric names, conditions and alerts, and any other information.

[0269] FIG. 28 illustrates a user interface display for a map view of asset tree data with a timeline. Careful consideration of interface 3900 of FIG. 28 will reveal many similarities to interface 3800 of FIG. 27. Noteworthy differences include the representation of mobile rather than stationary assets by the icons of interface 3900, and the presence of timeline control and navigation block 3940. Timeline control and navigation block 3940 in many respects is a timeline player for the map. A linear timeline is depicted by baseline 3944 and the end-to-end duration, and the start and end times are determined by timeframe component 3942 which in one embodiment may be implemented as a drop down selection list. Time position indicator 3946 represents the point in time within the time frame represented by baseline 3944 that is currently represented by the map view. As the time position indicator 3946 moves to different positions along baseline 3944 one can expect that asset icons for mobile assets such as asset icons 3920, 3910, and 3924, which may represent trucks in a fleet, for example, will move to different positions in the map so as to show their location of record at the time indicated by time position indicator 3946 in the timeframe. Geographic points of interest and constructs thereof, for example, Route 3910, may also be included in the map view. In one embodiment, geographic points of interest and constructs may be used in the determination of conditions and alerts. For example, a critical urgency level may be indicated where a truck asset departs more than 2 miles from the route asset. Time position indicator 3946 may be moved along the timeline by click and drag interactions, in one embodiment. Time position indicator 3946 may be moved along the timeline in a steady progressive way by activating play button 3948. Timeline control and navigation block 3940 also includes

timeline graph area **3950**. An embodiment may enable the graphing of one or more time series into timeline graph area **3950**. The data visualization in graph area **3950** (which may be in alignment with time baseline **3944**) may provide a quick visual clue as to time periods that may present more useful information. These and other embodiments are possible.

3.4 Asset Hierarchy Asset Groups

[0270] Aspects of embodiments heretofore described may be advantageously implemented with subject matter next discussed to provide novel embodiments related to the implementation of asset groups for an asset hierarchy. In the above-described embodiments, assets may be grouped that are similar in terms of sensors, capabilities, or other properties such as, but not limited to, location, seasonality, manufacturer make/model etc. These assets may be performing the same (or similar) tasks and measuring the same (or similar) metrics or related to each other based on the user's preference. In conventional systems, metrics, monitoring views, alerts, thresholds, and so on, are created for each asset individually. This can be a time and resource-intensive task consuming significant processing resources.

[0271] Embodiments described herein relate to an asset group interface driven by search-derived asset tree hierarchy. Embodiments described herein further relate to a control interface for metric definition specification for assets and asset groups that is driven by a search-derived asset tree hierarchy. An asset group associates multiple similar assets (also referred to herein as asset nodes) of the asset hierarchy at a group level. The asset group allows for the creation of group metrics, monitoring view, alerts, thresholds, and so on, for the asset group, which can be replicated to all or a subset of the assets that are members of the asset group. As a result, the time and resources utilized to generate the group metrics, monitoring view, alerts, thresholds, and so on, may be consumed once at the asset group level and then replicated to all or a subset of the assets that are members of the asset group, as compared to generating such metrics, views, alerts, thresholds, etc., in the conventional systems. Accordingly, a technical advantage of reduced processing resources and cycles is realized by implementing embodiments described herein. An appreciation for the solutions described above may be further developed by consideration of methods and user interface examples illustrated and discussed in relation to figures that follow.

[0272] FIG. 29 illustrates methods of an asset hierarchy monitoring and reporting system in one embodiment. The illustrated methods may represent a significant portion of the processing workflow to establish and operate an asset monitoring and reporting system (AMRS). Flowchart **4000** of FIG. 29 may beneficially be discussed in terms of an AMRS after the fashion illustrated in FIG. 10, such as an AMRS making combined use of functioning described for the data intake and query system **2120** and the asset system **2140** of FIG. 10. Furthermore, careful consideration of flowchart **4000** of FIG. 29 reveals many similarities to flowchart interface **2300** of FIG. 12. Noteworthy differences include the blocks **2342**, **2344**, **2346** and **2350** of flowchart **2300** combined into block **4012** of flowchart **4000**, and the presence of blocks **4014**, **4016**, **4018**, **4020**, and **4022**, which are discussed in further detail below. The description of flowchart **2300** of FIG. 12 is the same as described for identically numbered elements appearing in, and described

in relation to, depictions of user interface displays in earlier figures. The description of blocks **2342**, **2344**, **2346** and **2350** is similarly applied to block **4012** of flowchart **4000**.

[0273] At block **4014** of FIG. 29, configurations for asset groups may be created or edited. If one or more asset groups are being created, block **2342** may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation of new asset groups by the processing of block **4014**. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If the configuration of one or more asset groups is being edited, information of the existing one or more asset groups in CCC data store **2322** may be directly or indirectly represented in a display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block **4014** to construct a proper representation of one or more new or updated asset groups, and a representation of the new and/or updated asset groups may be reflected in the information of CCC data store **2322**. User interface displays of the processing of block **4014** may be caused to be displayed on a user interface device such as user computer **2314**, for example. Embodiments may vary as to the minimum and total amount of information that may be included for each asset group configuration. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of an asset group configuration, or any part or portion thereof, in CCC data store **2322**, in working storage during the processing of block **2342**, or elsewhere and at other times. The processing of block **2342** may be conducted, for example, by CCC console functions **2134** associated with an asset system **2140** as shown in FIG. 10. At the conclusion of the processing of block **4014** of FIG. 29, in an embodiment, information representing one or more new, changed, and/or unchanged asset group configurations may be found in CCC data store **2322**.

[0274] In an embodiment, the processing of block **2014** may include processing to create or edit definitions, specifications, indications, or the like, for one or more associations between and among one or more asset groups, on the one hand, and one or more assets represented in an asset hierarchy, on the other hand. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like, that are or may be used for the representation of such associations in computer storage at any place and time. In an embodiment, defined associations are reflected in the information of CCC data store **2322**, and here, as elsewhere, the representation of the defined associations in computer storage may be direct or indirect, expressed or implied, or otherwise.

[0275] At block **4016** of FIG. 29, definitions for asset group metrics (also referred to herein as "group metrics") may be created or edited. If one or more group metric definitions for asset group(s) are being created, block **4016** may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation of new group metrics definitions by the processing of block **4016**. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If the definition of one

or more group metrics is being edited, information of the existing one or more group metrics in CCC data store 2322 may be directly or indirectly represented in a display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 4016 to construct a proper representation of one or more new or updated group metric definitions, and a representation of the new and/or updated group metric definitions may be reflected in the information of CCC data store 2322. User interface displays of the processing of block 4016 may be caused to be displayed on a user interface device such as user computer 2314, for example. Embodiments may vary as to the minimum and total amount of information that may be included for each group metric definition. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of a group metric definition, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 4016, or elsewhere and at other times. The processing of block 4016 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as shown in FIG. 10. At the conclusion of the processing of block 4016 of FIG. 29, in an embodiment, information representing one or more new, changed, and/or unchanged group metric definitions may be found in CCC data store 2322.

[0276] In an embodiment, the processing of block 4016 may include processing to create or edit definitions, specifications, indications, or the like, for one or more associations between and among one or more group metrics, on the one hand, and one or more assets and one or more asset groups represented in an asset hierarchy, on the other hand. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like, that are or may be used for the representation of such associations in computer storage at any place and time. In an embodiment, defined associations are reflected in the information of CCC data store 2322, and here, as elsewhere, the representation of the defined associations in computer storage may be direct or indirect, expressed or implied, or otherwise.

[0277] At block 4018 of FIG. 29, definitions for conditions and/or alerts and/or actions corresponding an asset group may be created or edited. If one or more conditions and/or alerts and/or actions are being created, block 4018 may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation by the processing of block 4018 of new condition and/or alert and/or action definitions for an asset group. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If the definition of one or more asset group conditions and/or alerts and/or actions is being edited, information of the existing one or more asset group conditions/alerts/actions in CCC data store 2322 may be directly or indirectly represented in display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 4018 to construct a proper representation of the one or more new or updated asset group condition/alert/action definitions, and a representation of the new and/or updated asset group condition/alert/action definitions may be reflected in the infor-

mation of CCC data store 2322. User interface displays of the processing of block 4018 may be caused to be displayed on a user interface device such as user computer 2314, for example. Embodiments may vary as to the minimum and total amount of information that may be included for each asset group condition/alert/action definition. Embodiments may also vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of an asset group condition/alert/action definition, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 4018, or elsewhere and at other times. The processing of block 4018 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as shown in FIG. 10. At the conclusion of the processing of block 4018 of FIG. 29, in an embodiment, information representing one or more new, changed, and/or unchanged asset group condition/alert/action definitions may be found in CCC data store 2322.

[0278] In an embodiment, the processing of block 4018 may include processing to create or edit asset group definitions, specifications, indications, or the like, for one or more associations between and among one or more asset group conditions/alerts/actions, on the one hand, and one or more other defined objects, elements, or constructs, on the other hand. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like, that are or may be used for the representation of such associations in computer storage at any place and time. In an embodiment, defined associations are reflected in the information of CCC data store 2322, and here as elsewhere, the representation of the defined associations in computer storage may be direct or indirect, expressed or implied, and otherwise.

[0279] At block 4020 of FIG. 29, command, control, and configuration (CCC) information for the monitoring and/or reporting processing of the AMRS, for example with respect to asset groups and asset group metrics, as may appear in CCC data store 2322, may be created or edited. If the information is being created, block 4020 may present the user with an interface for specifying or indicating one or more items, aspects, characteristics, options, selections, or such for the creation of new CCC information by the processing of block 4020. In an embodiment, such an interface may include default values or selections for some or all of the information the interface enables a user to indicate or supply. If CCC information is being edited, existing information of CCC data store 2322 may be directly or indirectly represented in a display to a user via an interactive interface. The user may provide inputs via the user interface. Such user inputs may be used by the processing of block 4020 to construct a proper representation of the information and to reflect it in CCC data store 2322. User interface displays of the processing of block 4020 may be caused to be displayed on a user interface device such as user computer 2314, for example. Embodiments may vary as to the type and number of data items, formats, organizations, structures, representations, and the like that are or may be used for the storage of monitoring and/or reporting CCC information, or any part or portion thereof, in CCC data store 2322, in working storage during the processing of block 4020, or elsewhere and at other times. The processing of block 4020 may be conducted, for example, by CCC console functions 2134 associated with an asset system 2140 as

shown in FIG. 10. At the conclusion of the processing of block 4020 of FIG. 29, in an embodiment, information representing new, changed, and/or unchanged CCC information that determines, conditions, or otherwise influences the operation of monitoring and/or reporting aspects of the processing of an AMRS may be found in CCC data store 2322.

[0280] It is noted that as the illustrative example, method 4000 of FIG. 29 envisions an AMRS combining DIQ and asset system functionality as described in relation to FIG. 10. Furthermore, with respect to FIG. 29, CCC data store 2322 of FIG. 12 can find a counterpart in a shared CCC data store 2132 as described in relation to FIG. 10. Similarly, any discussion here in reference to CCC console processor 2134 of FIG. 10 embraces an embodiment of a shared CCC console as discussed above.

[0281] At block 4022 of FIG. 29, ongoing automatic and/or on-demand monitoring and/or reporting for an asset hierarchy, including asset groups and asset group metrics, as may be provided by an AMRS is conducted. The processing of block 4022 may utilize information reflected in CCC data store 2322 as described for the processing of other blocks of 4000, to determine, direct, condition, or otherwise influence its operational activity. In an embodiment, asset hierarchy monitoring activity may include data intake, internal data generation, computer-to-computer data transmission/presentation, and computer-to-person data presentation, for example, with a possible emphasis on data intake and generation aspects. In an embodiment, asset hierarchy reporting activity may include one or more of the same with a possible emphasis on data presentation aspects. The processing of block 4022 may be conducted, for example, by a combination of DIQ 2120 and asset system 2140 of FIG. 10, in an embodiment.

[0282] Examples for the processing of block 4022 for illustrative implementations, applications, or instantiations of an AMRS, follow. In one example, the processing of block 4022 may include monitoring sensors on a connected soldier in the battlefield and generate appropriate alerts for received impacts or heartbeat abnormalities. In one example, the processing of block 4022 may include monitoring process control equipment in a beverage processing facility and generate alerts for out-of-range temperatures or vibration anomalies. In one example, the processing of block 4022 may include monitoring activity in a workorder management system and may report impacts caused by scheduled downtime. In one example, the processing of block 4022 may include monitoring changes in process control settings, such as PID control settings, and report impacts on the accuracy of control loops. These are but a few illustrative examples.

[0283] An appreciation for the methods 4000 of FIG. 29 may be further developed by consideration of methods and user interface examples illustrated and discussed in relation to figures that follow.

[0284] FIGS. 30A and 30B illustrate methods for generation and management of asset groups and asset group metrics in view of an asset tree representation in control storage. Method 4100 of FIG. 30A illustrates a method as might be employed in an embodiment during the processing of block 4014 of FIG. 29, for example. Processing of block 4110 of FIG. 30A causes the display of a user interface on a user interface device. The user interface may be interactive enabling a user to both receive information from the AMRS (e.g., its CCC console processor) and to provide information

to the AMRS as enabled by the user interface. At block 4112, the AMRS receives an indication of asset nodes for inclusion in an asset group, which may be provided by user interaction with the user interface. In an embodiment, the asset node indications may represent some or all of the asset nodes for inclusion in an asset group of an asset hierarchy. In an embodiment, the asset node indications may be provided by user interaction in one or more forms, for example, selections of checkboxes associated with a particular asset nodes, segments of text in the form of a query language that specify asset nodes, selection of assets nodes from a drop-down menu, or others. At block 4114, indications of user group metric definitions for the asset group are received and processed. The processing of block 4114, in an embodiment, may involve receiving and processing an indication of a group metric definition that may be used to configure and/or define a metric definition that can be replicated to all assets that indicated as members of the asset group. The processing of block 4114 in an embodiment may receive and process an indication of a group metric definition that may be used to provide a group metric corresponding to an asset group identified in block 4112. In an embodiment, the user indications received and processed at block 4114 may be indications of the identification of input fields, calculation fields, time value fields, and so on provided in the user interface. Further detail on generating a group metric is discussed with respect to method 4150 of FIG. 30B below. The processing of block 4116, in an embodiment, may conclude with a representation of the determined asset group recorded in computer storage, perhaps in computer storage of the local working context, or perhaps in persistent computer storage of a CCC data store. These and other embodiments are possible.

[0285] Method 4150 of FIG. 30B illustrates a method as might be employed in an embodiment during the processing of block 4016 of FIG. 29, for example. Processing of block 4160 of FIG. 30B causes the display of a user interface on a user interface device. The user interface may be interactive enabling a user to both receive information from the AMRS (e.g., its CCC console processor) and to provide information to the AMRS as enabled by the user interface. At block 4162, the AMRS receives an indication of a metric definition specification for a metric of an asset of an asset group hierarchy, which may be provided by user interaction with the user interface. In an embodiment, the metric definition specification indications may represent some or all of the input values, calculation operations, and time factors for inclusion in a metric definition for an asset of an asset hierarchy. In an embodiment, the metric definition specification indications may be provided by user interaction in one or more forms such as, for example, selections of checkboxes associated with a particular assets, input values, calculations, time factors, operations, and so on. The metric definition indications may be provided by user interaction in one or more forms such as, for example, segments of text in the form of a query language that specify input values, calculations, time factors, operations, and so on. The metric definition indications may be provided by user interaction in one or more forms such as, for example, selection from a drop-down menu of input values, calculations, time factors, operations, and so on. Other embodiments are also possible.

[0286] At block 4164, indications of user group metric definitions for the asset group are received and processed. The processing of block 4164, in an embodiment, the AMRS

receives an indication of a group metric definition specification for a group metric of an asset group of an asset group hierarchy, which may be provided by user interaction with the user interface. In an embodiment, the group metric definition specification indications may represent some or all of the input values, calculation operations, and time factors for inclusion in a group metric definition for an asset of an asset hierarchy. In an embodiment, the group metric definition specification indications may be provided by user interaction in one or more forms such as, for example, selections of checkboxes associated with a particular group metrics of assets in the asset group, input values, calculations, time factors, operations, and so on. The group metric definition indications may be provided by user interaction in one or more forms such as, for example, segments of text in the form of a query language that specify input values, calculations, time factors, operations, and so on. The group metric definition indications may be provided by user interaction in one or more forms, such as, for example, selection from a drop-down menu or otherwise of input values, calculations, time factors, operations, and so on. Other embodiments are also possible. In an embodiment, the group metric definition is replicated to all assets that are members of the asset group. The processing of block 4166, in an embodiment, may conclude with a representation of the determined metric definition and group metric definition recorded in computer storage, perhaps in computer storage of the local working context, or perhaps in persistent computer storage of a CCC data store. These and other embodiments are possible.

[0287] An appreciation for the methods 4100 and 4150 of FIGS. 30A and 30B may be further developed by consideration of methods and user interface examples illustrated and discussed in relation to figures that follow.

3.4.1 Creating Asset Groups

[0288] The discussion that follows relates to the generation and configuration of asset groups performed by the AMRS machinery. Figures that follow relate largely to interfaces that may be utilized by the AMRS during the processing described and contemplated for block 4014 of FIG. 29 as performed by a monitor/reporter processor such as 2142 of FIG. 10, for example.

[0289] FIG. 31 illustrates a user interface display for an asset group console function. User interface display 4200 is such as might be caused to display during the processing of block 4014 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. User interface display 4200 of FIG. 31 is shown to include system header bar 2402, application information and menu bar 2902, asset hierarchy display area 2920, and asset overview display area 2940. System header bar 2402 is as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Application information and menu bar 2902 is comparable to other application information and menu bars depicted and described in relation to earlier appearing figures, such as bar 2602 of FIG. 17. Application menu selectors 4212, 4214, 4216, 4218 may include a monitor selector 4212, a browse selector 4214, an analysis selector 4216, and a settings selector 4218. Browse selector 4214 is depicted with a different background color than that of the other selectors of

interface 4200, indicating a default selection of the browse selector 4214 for user interaction with the functions of browse selector 4214.

[0290] Asset hierarchy display area 4220 may enable a user to identify one or more assets of an asset tree for which asset group information may be created, edited, deleted, or otherwise processed. Asset overview display area 4230 may enable a user to view and interact with information and processing for a particular asset of the asset hierarchy.

[0291] Asset hierarchy display area 4220 is shown to include asset tab 4222 for “Assets”, groups tab 4223 for “Groups”, asset filter component 4224, and asset node list 4226. In asset hierarchy display area 4220 of interface 4200, asset tab 4222 is depicted with an underline indicating a default selection of the asset tab 4222 associated with display and browsing of asset nodes of an asset hierarchy.

[0292] Asset node list 4226 includes a node list entry for each of one or more of the nodes in an asset hierarchy. In one embodiment, initially and by default, node list 4226 includes a node list entry for every node in the asset hierarchy except for the root node. In an embodiment, a user may interact with asset filter component 4224 of interface 4200 to indicate to the AMRS filter criteria as may be applied to the nodes of the asset hierarchy before populating list 4226. For example, a user may enter the word “s_engine” into a text box of filter component 4224 and the AMRS upon receiving that indication may update asset tree node list 4226 to include entries for only those nodes of the asset tree whose names include the word “s_engine.” In such an example, the displayed asset tree node list 4226 would be shortened to show only the “100_1100_s_engine” node entry, the “100_1101_s_engine” node entry, and the “100_1102_s_engine” node entry. The display of asset tree node list 4226 may depict the hierarchical relationship among the nodes using an indentation scheme such as that described in regard to the node list display of 2810 of FIG. 17. Asset tree node list 4226 of FIG. 31 may include interactive expansion/contraction selectors (e.g., downward-facing carrot) for one or more node entries to enable a user to expand and/or contract sections of the asset tree hierarchy for ease of viewing and/or navigation.

[0293] Asset overview display area 4230 is shown to include area title “Asset” and an identification 4232 of the asset node selected in the asset hierarchy display area 4220. For example, identification 4232 identifies the asset node of interest as “100_Series.100_1100.s_engine,” which is asset node “100_1100_s_engine” 4228 of asset tree node list 4226 of interface 4200. Asset node “100_1100_s_engine” 4228 is depicted with a different background color than that of the other entries in asset tree node list 4226 of interface 4200, indicating a default selection of the asset node associated with that entry or indicating a prior user interaction with that entry to effect such a selection. The selection of an entry of asset tree node list 4226 may result in the display of information and user interface elements in the asset overview display 4230 that are pertinent to the asset represented by the selected asset node, such as selected asset node 4228.

[0294] Asset overview display area 4230 also includes a metrics tab 4234. In asset overview display area 4230 of interface 4200, metrics tab 4234 is depicted with an underline indicating a default selection of the metrics tab 4234 associated with display and browsing of metrics 4236 of a particular asset node (e.g., asset node 4228) of an asset hierarchy 4226. As shown in interface 4200, asset node 4228

(“100_Series.100_1100.s_engine”) is associated with 4 metrics **4236**, including “fuel_consumption”, “oil_level”, “rpm”, and “temperature”. These metrics **4236** are displayed in a tabular format in interface **4200**, may also be displayed in different views. Each metric **4236** may have a corresponding type **4240** (e.g., “raw” data collected from sensors or logs corresponding to the asset, calculated, group reference, etc.) and action **4242** (e.g., analyze, delete).

[0295] In an embodiment, metrics list **4236** may be populated by the AMRS with an entry for each already-defined metric known to the AMRS. In an embodiment, metrics list **4236** may be populated by the AMRS with an entry for each already-defined metric known to the AMRS and associated with at least one node of the asset tree hierarchy represented in whole or in part in **4220**. In an embodiment, metrics list **4236** may be populated by the AMRS with an entry for one or more metrics known to the AMRS to be associated with any manufacturer (or other metadata information item) that is associated with any asset node of the asset tree hierarchy represented in whole or in part in **4220**. These and other embodiments are possible. In an embodiment, population of metrics list **4236** by the AMRS may be influenced by prior user interaction with metrics search component **4244** to supply a filter criteria.

[0296] Individual entries of metrics list **4236** may each be enabled for user interaction so as to enable a user to indicate a selection of one of the entries of the list, in an embodiment. The selection of an entry of metrics list **4236** may result in the display of information and user interface elements in another user interface (not shown). User interaction with “Add new Metric” action button **4238** may be indicated to the AMRS, which may in response present a modified or alternate user interface display enabling a user to indicate sufficient information to define a new metric.

[0297] A user interaction with “Add to Group” action button **4250** that is indicated to the AMRS may cause the AMRS to engage processing to effect the addition of the currently-selected asset node **4232** to an asset group. In one embodiment, in response to an indication of user interaction with “Add to Group” action button **4250** of FIG. 31, AMRS may engage processing to effect the addition of one or more nodes to an asset group by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define an asset group for the asset node to join in the asset tree hierarchy. In one embodiment, the AMRS may present a user interface similar to interface **4900** of FIG. 38 that is enabled for data entry by the user such that the user can manually add an asset node to an asset group of the asset node hierarchy. These and other embodiments are possible.

[0298] FIG. 32 illustrates a user interface display for an asset groups console function for adding a new asset group. User interface display **4300** is such as might be caused to display during the processing of block **4014** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. Interface **4300** FIG. 32 is shown to include system header bar **2402**, application information and menu bar **4202**, application menu selectors **4212-4218**, assets tab **4222**, groups tab **4223**, asset group action buttons **4302**, **4304**, and active group display area **4306**. System header bar **2402**, application information and menu bar **4202**, and application menu selectors **4212-4218** are as described for identically numbered elements appearing in, and described

in relation to, depictions of user interface displays in earlier figures. Similarly, assets tab **4222** and groups tab **4223** are as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures.

[0299] In FIG. 32, the groups tab **4223** is depicted with an underline indicating a default selection of the groups tab **4223** associated with display and browsing of a particular asset group of an asset hierarchy. Tab interfaces are understood in the art and will not be elaborated here. When groups tab **4223** is the active tab selection, asset group action buttons **4302**, **4304** are shown in interface **4300**. Asset group action button can be shown when there are no asset groups created and/or available for selection in the AMRS. Similarly, active group display area **4306** displays a message, such as “no group selected”, when there are no groups available and/or selected in the groups tab **4223**.

[0300] Asset group action buttons **4304** is a default asset group action button persistently displayed whenever the groups tab **4223** is the active tab. In one embodiment, in response to an indication of user interaction with either asset group action button **4302**, **4304** of FIG. 32, AMRS may engage processing to effect the addition of one or more asset groups by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define a new asset group for the asset tree hierarchy.

[0301] FIG. 33 illustrates a user interface display for an asset group creation console function for specifying an asset group name. User interface display **4400** is such as might be caused to display during the processing of block **4014** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **4400** is the same as interface **4300** of FIG. 32, with the addition of asset group name window **4402** displayed in response to user selection of asset group action buttons **4302**. Asset group name window **4402** includes an asset group name input box **4404** and add action button **4406**. Asset group name input box **4404** enables a user to input a name of a new asset group to be created. Selection of the add action button **4406** causes a new asset group with a name identified in asset group name input box **4404** to be created.

[0302] FIG. 34 illustrates a user interface display for an asset group creation console function for adding a new asset to an asset group. User interface display **4500** is such as might be caused to display during the processing of block **4014** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10.

[0303] User interface display **4500** is the same as interface **4400** of FIG. 33. Interface **4500** illustrates the groups tab **4223** after the add action button **4406** in interface **4400** of FIG. 33 has been selected. Groups tab **4223** of interface **4500** displays the new asset group “siemens engine” **4502** as the selected active asset group. Asset group “siemens engine” **4504** is also displayed in asset group detail area **4508**. When there are no assets yet assigned to an asset group, such as with the newly-created “siemens engine” asset group **4502**, asset group detail area **4508** can display a message the “You don’t have any assets in the group” and can include an add asset action button **4506**. In one embodiment, in response to an indication of user interaction with the add asset action button **4506** of FIG. 34, AMRS may engage processing to effect the addition of one or more

assets to the asset group by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define a new asset for the asset group.

[0304] FIG. 35 illustrates a user interface display an asset group console function for specifying assets to add to an asset group. User interface display 4600 is such as might be caused to display during the processing of block 4014 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. User interface display 4600 displays an add asset overlay window 4602 on top of interface 4500 of FIG. 34. The add asset overlay window 4602 of interface 4600 is a user interface to add assets to an asset group that is presented add asset action button 4506 of interface 4500 of FIG. 34 has been selected.

[0305] Add asset overlay window 4602 enables the addition of one or more assets to an asset group. Add asset overlay window 4602 includes an asset node list 4604 and asset node information 4606. Asset node list 4606 includes a node list entry for each of one or more of the nodes in an asset hierarchy. Asset node information 4606 may identify a total number of asset nodes to select from and a current number of selected asset nodes in the asset node list 4604.

[0306] In one embodiment, initially and by default, asset node list 4604 includes a node list entry for every node in the asset hierarchy except for the root node. In an embodiment, a user may interact with asset search component 4610 of interface 4600 to indicate to the AMRS filter criteria as may be applied to the nodes of the asset hierarchy before populating list 4604. For example, a user may enter the word “s_engine” into a text box of search component 4610 and the AMRS upon receiving that indication may update asset tree node list 4604 to include entries for only those nodes of the asset tree whose names include the word “s_engine.” In such an example, the displayed asset tree node list 4604 would be shortened to show only the “100_series.100_1100.s_engine” node entry, “100_series.100_1101.s_engine” node entry, and the “100_series.100_1102.s_engine” node entry. The display of asset node list 4604 may depict the hierarchical relationship among the nodes using an indentation scheme such as that described in regard to the node list display of 2810 of FIG. 17. Asset node list 4604 of FIG. 35 may include an interactive check box for one or more node entries to enable a user to identify a selection of one or more asset nodes through interaction with the checkboxes. For example, a user may interact with a number of checkboxes to place them in a selected state, and the set of asset node entries with selected checkboxes may be used to associate the asset tree nodes represented by those entries with a metric being created or edited.

[0307] Add asset overlay window 4602 may also include a current view selector 4608 that allows a user to change the set of asset nodes displayed in the asset node list 4604. Current view selector 4608 is shown in a “Show All Assets” state resulting in all asset nodes in the asset tree hierarchy being displayed in the asset node list 4604. In some implementations, current view selector 4608 may be modified by a user to cause a subset of asset nodes of the asset tree hierarchy to be displayed in the asset node list 4604.

[0308] As shown in FIG. 35, asset nodes “100_series.100_1100.s_engine”, “100_series.100_1101.s_engine”, and the “100_series.100_1102.s_engine” are currently selected (via checkboxes in asset node list 4604) for inclusion in the asset group. Selection of the add action button 4614 may engage

processing to effect the addition of one or more asset nodes to the asset group. Conversely, selection of the cancel action button 4612 may engage processing to effect the cancellation of the processing of the add asset overlay window 4602 with a result of no asset nodes being added to the asset group.

[0309] FIG. 36 illustrates a user interface display for an asset group members console. User interface display 4700 is such as might be caused to display during the processing of block 4014 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 4700 illustrates the groups tab 4223 after the add action button 4614 in interface 4600 of FIG. 35 has been selected. Groups tab 4223 of interface 4700 displays the asset group “siemens engine” 4502 as the selected active asset group. Asset group “siemens engine” 4504 is also displayed in asset group detail area 4508. Asset group detail area 4508 further includes a members tab and a group metrics tab 4704. Members tab 4702 details the asset node members of an asset group. Group metric tab 4704 details metrics configurations for the asset group, which is discussed in more detail further below. The add view action element 4706 engages processing to generate a new view for the asset group.

[0310] Interface 4700 illustrates that members tab 4702 is the active tab. The members tab 4702, when activated, provides an asset node member list 4710 identifying the asset nodes that have been added as members of the asset group. Individual entries of asset node member list 4710 may each be enabled for user interaction so as to enable a user to indicate a selection of one of the entries of the list, in an embodiment. The selection of an entry of asset node member list 4710 may result in the display of information and user interface elements in another user interface (not shown). In an embodiment, a user may interact with asset search component 4712 of interface 4700 to indicate to the AMRS filter criteria as may be applied to the nodes of the asset node member list 4710 to populate list 4710. Individual entries of asset node member list 4710 may each have an associated action interface element 4714. Different actions may be provided in action interface element 4714, such as effecting removal of the corresponding asset node from the asset group. Asset node list filter element 4716 may allow a user to indicate how many asset node members to display in the asset node member list 4710.

[0311] In one embodiment, in response to an indication of user interaction with the add asset member action button 4720 of FIG. 36, AMRS may engage processing to effect the addition of one or more asset nodes to the asset group by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define a new asset for the asset group. In another embodiment, in response to an indication of user interaction with the delete asset group action button 4730 of FIG. 36, AMRS may engage processing to effect the deletion of the asset group by presenting a user interface, such as a modal window, that enables a user to specify information to cause the asset group, in its entirety, to be deleted.

[0312] In some embodiments, a new asset group may be automatically created based on rule-based criteria. For example, identification of assets to include in an asset group may be performed by receiving, into memory, criteria associated with membership in the asset group based at least in part on user input. Subsequently, one or more asset nodes of

the asset hierarchy comporting with said criteria are identified to associate with the asset group.

[0313] Referring back to FIG. 31, a new asset group can also be created in response to an indication of user interaction with the “Add to Group” action button 4250 of FIG. 31. Responsive to activation of “Add to Group” action button 4250, AMRS may engage processing to effect the addition of the active asset node (e.g., asset node 4228, 4232) detailed in asset overview display 4230 to an asset group by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to add the asset to the asset group.

[0314] FIG. 37 illustrates a user interface display for adding an in-focus asset to an asset group. User interface display 4800 is such as might be caused to display during the processing of block 4014 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 4800 is the same as interface 4200 of FIG. 31, with the addition of asset group name input box 4802 displayed in response to user selection of “Add to Group” action button 4250 of FIG. 31. Asset group name input box 4802 enables a user to input a name of an asset group (new or already existed) for the active asset node to be added. Asset group suggestion box 4804 may be provided to a user to indicate suggestion based on natural language processing (NLP) applied to the text input into the asset group name input box 4802 as the user inputs the text. In response to receiving an indication (e.g., selection of enter or input key by user) of final input of the asset group name, the active asset node is added as a member of the indicated asset group.

[0315] FIG. 38 illustrates a user interface display for addition of an asset member to an asset group. User interface display 4902 is such as might be caused to display during the processing of block 4014 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 4900 is the same as interface 4200 of FIG. 31, with the addition of asset group identifier element 4902 displayed in response to addition of the asset node to an asset group, such as via asset group name input box 4802 of FIG. 37.

[0316] FIG. 39 illustrates a user interface display for an asset group browsing console for addition of assets to an asset group. User interface display 5000 is such as might be caused to display during the processing of block 4014 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 5000 is similar to interface 4700 of FIG. 36. However, interface 5000 illustrates a single asset node member 5002 (“series_100.100_1100.s_engine”) as a member of the “siemens engine” asset group, with the addition of asset group identifier element 4902 displayed in response to addition of the asset node to an asset group, such as via asset group name input box 4802 of FIG. 37. In one embodiment, in response to an indication of user interaction with the add asset member action button 4720 of FIG. 39, AMRS may engage processing to effect the addition of one or more asset nodes to the asset group by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define a new asset for the asset group. In one embodiment, interface 4600 of FIG. 35 may be presented to the user to enable the user to specify information sufficient to define a new asset for the asset group.

[0317] The user interfaces already discussed in relation to asset group generation and configuration, namely, the user interfaces illustrated and discussed in relation to FIGS. 31-39, have largely related to interfaces employed by a command, control, and configuration console processor (such as CCC console 2134 of FIG. 10). These interfaces enable a user, such as a system administrator or operator, to manipulate the virtual levers, buttons, dials, and switches (embodied in the information of a CCC data store such as 2132 of FIG. 10) that control the operation of the asset monitoring and reporting machine. The focus now turns with the discussion of the figures that follows to generation and configuration of group metrics for an asset group performed by the AMRS machinery. Figures that follow relate largely to interfaces that may be utilized by the AMRS during the processing described and contemplated for block 4016 of FIG. 29 as performed by a monitor/reporter processor such as 2142 of FIG. 10, for example.

3.4.2 Creating Group Metrics for Asset Groups

[0318] FIG. 40 illustrates a user interface display for group metrics of an asset group console. User interface display 5100 is such as might be caused to display during the processing of block 4016 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 5100 is similar to interface 4700 of FIG. 36. Interface 5100 illustrates the group metrics tab 4704 for an asset group. Group metrics tab 4704 details metrics configured for an asset group 4502, 4504, selected as an active asset group in UI 5100.

[0319] Interface 5100 illustrates that group metrics tab 4704 as the active tab. The group metrics tab 4704, when activated, provides a group metrics overview display 5102 portion of interface 5100. The group metrics overview display 5102 includes a group metrics list 5104 identifying the metrics that correspond to the selected asset group 4502. A group metric is a metric with the same name that can be available across all asset node members of an asset group. For example, Asset A has raw metrics (e.g., metrics provided from a customer’s source system) Cat, Dog. Asset B has raw metrics Car, Bus. Asset C has raw metrics Carrot, Apple. An asset group can be created for Assets A, B and C, even though they do not share any similar metrics. Also, the user could decide that he or she would like to create an asset group with assets that has metrics that starts with the letter C for example (exaggerating here to show flexibility of group). Consequently, the group metrics that are available after an asset group is created include the union of all metrics on the assets. An asset group is designed in this matter because assets of an asset group may not have related/similar metrics in all cases. This approach provides more flexibility to the user to create an asset group based on the user’s preference.

[0320] Interface 5100 indicates that there are 4 metrics in the group metrics list 5104 associated with the “siemens engine” asset group, where the metrics are “fuel_consumption”, “oil_level”, “rpm”, and “temperature”. Individual entries of group metrics list 5104 may each be enabled for user interaction so as to enable a user to indicate a selection of one of the entries of the list, in an embodiment. The selection of an entry of group metrics list 5104 may result in the display of information and user interface elements in another user interface.

[0321] FIG. 41 illustrates a user interface display for renaming a group metric of an asset group. User interface display 5200 is the same as interface 5100 of FIG. 40, with the addition of rename group metric window 5202 displayed in response to user selection of an individual group metric (e.g., “fuel_consumption”) of group metrics list 5104 in interface 5100. Rename group metric window 5202 includes a group metric name box 5204 and a group metric description box 5206. Boxes 5204 and 5206 are pre-filled with the pertinent identifying information corresponding to the selected group metric. Boxes 5204 and 5206 are editable to allow a user to make changes to the pre-filled information in boxes 5204, 5206. Selection of the cancel button 5210 causes the rename group metric window 5202 to be closed without any changes saved to the corresponding group metric. Selection of the rename button 5220 causes the corresponding group metric to be updated with the information in the group metric name box 5204 and group metric description box 5206 and saved.

[0322] Referring back to FIG. 40, in an embodiment, each group metric in the group metrics list 5104 has a corresponding type 5106 and action 5108. The group metric type 5106 may be, but is not limited to, a “member reference” or “calculated” type. “Member reference” may refer to a raw metric associated with an individual member and applied to all members in the asset group. A calculated type may refer to a group metric configured by a user for the group. “Calculated” metrics may be generated using a formula builder interface that is described in further detail below. An Action 5108 associated with a group metric may include, but is not limited to, analyze or delete. In one embodiment, in response to an indication of user interaction with the analyze action 5108 of FIG. 40, AMRS may engage processing to effect the monitoring and reporting on the group metric by presenting a user interface, such as a modal window, that enables a user to configure and view monitoring and reporting on the group metric. In one embodiment, in response to an indication of user interaction with a delete (not shown) action 5108 of FIG. 40, AMRS may engage processing to effect delete the particular group metric from the asset group. In one embodiment, calculated-type group metrics may have a corresponding delete action 5108, while member reference-type group metrics do not have a corresponding delete action 5108.

[0323] In one embodiment, in response to an indication of user interaction with an add group metric action button 5110 of FIG. 40, AMRS may engage processing to effect the addition of a group metric to the asset group by presenting a user interface, such as a modal window, that enables a user to specify information sufficient to define a new group metric for the asset group.

[0324] FIG. 42 illustrates a user interface display for adding a group metric to an asset group. User interface display 5300 is the same as interface 5100 of FIG. 40, with the addition of add group metric window 5302 displayed in response to user selection of add group metric action button 5110 in interface 5100. Add group metric window 5302 includes a group metric name box 5304 and a group metric description box 5306. Boxes 5304 and 5306 are initially presented to the user as blank input boxes. Boxes 5204 and 5206 are editable to allow a user to input information regarding the new group metric in boxes 5204, 5206. As illustrated in interface 5300, a new group metric “fuel_consumption_AVG” is provided. Selection of the cancel

button 5310 causes the add group metric window 5302 to be closed without any changes saved for the new group metric. Selection of the add button 5320 causes the new group metric to be added to the asset group with the information indicated in the group metric name box 5304 and group metric description box 5306.

3.4.2.1 Formula Builder for Creating Group Metrics for an Asset Group

[0325] FIG. 43 illustrates a user interface display for configuring a group metric of an asset group console. User interface display 5400 is such as might be caused to display during the processing of block 4016 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 5400 may also be referred to as a formula builder, and may be utilized in embodiments to configure and generate custom metrics and group metrics for assets and asset groups of an asset hierarchy. In one embodiment, interface 5400 is displayed in response to selection of the add new metric button 5110 of interface 5100 of FIG. 40. User interface display 5400 of FIG. 43 is shown to include system header bar 2402, application information and menu bar 5402, asset hierarchy display area 5404, group metric determination area 5406, and group metric function area 5408. System header bar 2402 is as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Application information and menu bar 5402 is comparable to other application information and menu bars depicted and described in relation to earlier appearing figures, such as bar 2602 of FIG. 17. Application information and menu bar 5402 may also include an identification of the group metric (e.g., “siemens engine.fuel_consumption_AVG”) that is being configured in interface 5400.

[0326] Asset hierarchy display area 5404 may enable a user to identify one or more assets of an asset tree for which a group metric may be configured, edited, deleted, or otherwise processed. Asset hierarchy display area 5404 is comparable to other asset hierarchy display areas depicted and described in relation to earlier appearing figures, such as display area 4220 of FIG. 31.

[0327] Asset hierarchy display area 5404 is shown to include asset tab 5410 for “Assets”, groups tab 5412 for “Groups”, other tab 5414, asset filter component 5418, and asset node list 5416. In asset hierarchy display area 5404 of interface 5400, asset tab 5410 is depicted with an underline indicating a default selection of the asset tab 5410 associated with display and browsing of asset nodes of an asset hierarchy.

[0328] Asset node list 5416 includes a node list entry for each of one or more of the nodes in an asset hierarchy. In one embodiment, initially and by default, asset node list 5416 includes a node list entry for every node in the asset hierarchy except for the root node. Asset node list 5416 also includes metrics associated with each asset. In an embodiment, a user may interact with asset filter component 5418 of interface 5400 to indicate to the AMRS filter criteria as may be applied to the nodes of the asset hierarchy before populating list 5416. For example, a user may enter the word “s_engine” into a text box of filter component 5418 and the AMRS upon receiving that indication may update asset tree node list 5416 to include entries for only those nodes of the asset tree whose names include the word “s_engine.” In such

an example, the displayed asset tree node list **5416** would be shortened to show only the “100_1100_s_engine” node entry, the “100_1101_s_engine” node entry, and the “100_1102_s_engine” node entry, and their associated asset metrics. The display of asset tree node list **5416** may depict the hierarchical relationship among the nodes using an indentation scheme such as that described in regard to the node list display of **2810** of FIG. 17. Asset tree node list **5416** of FIG. 43 may include interactive expansion/contraction selectors (e.g., downward-facing carrot) for one or more node entries to enable a user to expand and/or contract sections of the asset tree hierarchy for ease of viewing and/or navigation.

[0329] Group metric determination area **5406** may enable a user to configure, edit, view, delete, and otherwise interact with information and processing for a group metric of an asset group. Group metric determination area **5406** includes an editable input area **5422** to display or specify information of a group metric definition. Group metric determination area **5406** may be used to display or specify information of a group metric definition related to aspects of a search query that determine a value for the group metric from the data selected by the search query. The search query may be generated in the group metric determination area **5406** based on user interaction with interface **5400**, as described in further detail below.

[0330] Group metric determination area **5406** also includes a scheduling component **5424** and a clear action button **5426**. The scheduling component **5424** may be used to display or specify information of when and/or how often the group metric calculation entered into editable input area **5422** is performed. Scheduling component **5424** is shown as a drop-down selection list displaying 1 minute as the current selection by default, prior user interaction, or otherwise. In one embodiment, the drop-down selection list of scheduling component **5424** may include, 1 minute, 5 minutes, 10 minutes, 15 minutes, 20 minutes, 30 minutes, 1 hour, 2 hours, 4 hours, 6 hours, 12 hours, 1 day, 2 days, 3 days, and 7 days, as available options, for example. The clear action button **5426** enables a user to indicate to delete an input into the editable input area **5422**. The ARMS receiving such an indication may perform such a deletion and refresh or update the editable input area **5422**.

[0331] Group metric function area **5408** may be used to specify a calculation option to include in the editable input area **5422** of group metric determination area **5406**. Group metric function area **5408** is shown as multiple expandable menus of selectable calculation options **5430**. The multiple expandable menus may include, but are not limited to, conditional **5432**, mathematical **5434**, statistical **5436**, and trigonometry **5438**. Interface **5400** is illustrated with the statistical menu **5436** in an expanded state that displays a list **5440** of selectable statistical calculation options including avg(X) (average calculation operation), max(X) (maximum calculation operation), median (X) (median calculation operation), min(X) (minimum calculation operation), range (x) (range calculation operation), stdev(x) (standard deviation calculation operation), sum(x) (summation calculation operation), var(X) (variance calculation operation), varp(X) (variance population calculation operation). More or less selectable calculation options than shown may be provided in the statistical menu **5436**. Each of the conditional **5432**, mathematical **5434**, and trigonometry **5438** expandable menus may also include sets of selectable calculation options corresponding to the particular type of menu.

[0332] In one embodiment, in response to an indication of user interaction with any one of the selectable calculation options **5440** of FIG. 43, AMRS may engage processing to cause the selected calculation option to be provided as a calculation in the editable input area **5422** to define the group metric being configured in interface **5400**, as described in more detail below with respect to FIG. 44.

[0333] Selection of the Cancel button **5450** causes the interface **5400** for configuration of the group metric to be closed without any changes saved for the group metric of the asset group. Selection of the Save button **5460** causes the defined configuration in editable input area **5422** for the group metric to be saved for the asset group.

[0334] FIG. 44 illustrates a user interface display for configuring a group metric of an asset group console with editable group metric definitions based on asset metrics. User interface display **5500** is such as might be caused to display during the processing of block **4016** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **5500** is the same as interface **5400** of FIG. 43, with the addition of configuration input **5502** to define a group metric (e.g., “siemens.engine.fuel_consumption_AVG” identified in application information and menu bar **5404**) displayed in the editable input area **5422** of interface **5500**. Configuration input **5502** is displayed in response to user selection of various selectable elements of the interface **5500** to define the configuration of the group metric in the editable input area **5422** of interface **5500**.

[0335] In one embodiment, configuration input **5502** may be generated in response to selection of one of more selectable calculation options. As illustrated in interface **5500**, in response to user selection of calculation option **5506** for the function “avg (x)”, the calculation option avg (x) **5504** is added to the editable input area **5422**. In one embodiment, parameters of the selected calculation option **5504** may be added to the configuration input **5502** via user selection of asset metrics from the asset tab **5410**. As illustrated, in response to user selection of asset metrics **5520** (“100_series.1100.s_engine.fuel_consumption”), **5522** (“100_series.1101.s_engine.fuel_consumption”), **5524** (“100_series.1102.s_engine.fuel_consumption”) from the asset tab **5410**, corresponding asset metrics are added as corresponding parameters **5510** (“100_series.1100.s_engine.fuel_consumption”), **5512** (“100_series.1101.s_engine.fuel_consumption”), **5514** (“100_series.1102.s_engine.fuel_consumption”) to the calculation option **5504** in the configuration input **5502**.

[0336] FIG. 45 illustrates a user interface display for configuring a group metric of an asset group console with editable group metric definitions based on asset group metrics. User interface display **5600** is such as might be caused to display during the processing of block **4016** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **5600** is the same as interface **5400** of FIG. 43. However, configuration input **5602** in the editable input area **5422** is now generated for the same group metric (“siemens.engine.fuel_consumption_AVG” identified in application information and menu bar **5404**) as shown in interface **5500** using asset group metrics instead of individual asset metrics. Configuration input **5602** is displayed in response to user selection of

various selectable elements of the interface **5600** to define the configuration of the group metric in the editable input area **5422** of interface **5600**.

[0337] In one embodiment, configuration input **5602** may be generated in response to selection of one or more selectable calculation options, such as “avg(X)” similar to as described with respect to interface **5500**. In interface **5600**, parameters of the selected calculation option in configuration input **5602** may be added via user selection of asset group metrics from the group tab **5412**. As illustrated, in response to user selection of the asset group metric **5604** (“siemens.engine.fuel_consumption”) from the groups tab **5412**, a corresponding asset group metric is added as a parameter **5606** (“siemens.engine.fuel_consumption”) to the calculation option in the configuration input **5602**. This allows a user to generate a group metric that is replicated to all members of an asset group to improve utilization of computing and processor resources.

[0338] FIG. 46 illustrates a user interface display for an asset group console with group metrics browsing. User interface display **5700** is such as might be caused to display during the processing of block **4016** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **5700** is the same as interface **5100** of FIG. 40. In interface **5700**, group metrics list **5104** now identifies the new calculated group metric “fuel_consumption_AVG” **5702** that was configured and defined in interface **5600** of FIG. 45. In one embodiment, interface **5700** is displayed in response to a user selecting the Save button **5460** of interface **5600**. New calculated group metric **5702** is associated with a corresponding type **5106** of calculated **5074** and corresponding actions **5108** of analyze and delete **5706**.

[0339] FIG. 47 illustrates a user interface display for an asset group metric configuration console. User interface display **5800** is such as might be caused to display during the processing of block **4016** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **5800** is the same as interface **5400** of FIG. 43. In interface **5800**, asset metrics list **5416** now identifies the new calculated group metric “fuel_consumption_AVG” **5802** that was configured and defined in interface **5600** of FIG. 45. As illustrated, new calculated group metric is shown in each member of the asset group. As such, when the asset tab **5410** is the active tab, the new group metric **5802** may appear multiple times (e.g., with each asset member of the asset group). In one embodiment, interface **5800** is displayed in response to selection of the add new metric button **5110** of interface **5100** of FIG. 40. The new calculated group metric **5802** is shown as a selectable input for use as a parameter in editable input area **5422** when configuring and defining a group metric for an asset group.

[0340] FIG. 48 illustrates a user interface display for an asset group console with group metrics replicated to an asset. User interface display **5900** is such as might be caused to display during the processing of block **4016** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **5900** is the same as interface **4200** of FIG. 31 and as interface **4900** of FIG. 38. In interface **5900**, asset details for an individual asset node **4228** (“100_series.100_1100.s_engine”) selected

for viewing from assets tab **4222** are shown. The metrics **4236** of the asset node **4228** are shown when metrics tab **4234** is active. When a group metric is created for an asset group **4902** (e.g., “siemens engine”) to which asset node **4228** is a member, the group metric associated with the asset group is populated to all members of the asset group. As such, the metrics **4236** displayed for the individual asset node **4228** now identifies the calculated group metric “fuel_consumption_AVG” **5902** as an asset metric that was configured and defined for the asset node **4228**. Calculated group metric **5902** is associated with a corresponding type **4240** of “group reference” **5904**.

[0341] The user interfaces already discussed in relation to group metric generation and configuration for an asset group, namely, the user interfaces illustrated and discussed in relation to FIGS. 40-48, have largely related to interfaces employed by a command, control, and configuration console processor (such as CCC console **2134** of FIG. 10). These interfaces enable a user, such as a system administrator or operator, to manipulate the virtual levers, buttons, dials, and switches (embodied in the information of a CCC data store such as **2132** of FIG. 10) that control the operation of the asset monitoring and reporting machine. The focus now turns with the discussion of the figures that follows to asset group monitoring and reporting performed by the AMRS machinery. Figures that follow relate largely to interfaces that may be utilized by the AMRS during the processing described and contemplated for block **4020** of FIG. 29 as performed by a monitor/reporter processor such as **2142** of FIG. 10, for example.

3.4.3 Asset Group Monitoring/Reporting

[0342] Referring back to FIG. 36, in an embodiment, interface **4700** illustrated an asset group view for an asset group. Interface **4700** includes a groups tab **4223** that displays the asset group “siemens engine” **4502** as the selected active asset group. Along with a members tab **4702** and group metrics tab **4704** detailing asset node members and groups metrics corresponding to the particular group, an add view action element **4706** is provided. As discussed above, the add view action element **4706** engages processing to generate a new view for the asset group.

[0343] In one embodiment, a view may refer to a custom monitoring or reporting presentation for asset tree information and may be termed a display or presentation view. A view instance produced by the AMRS during the performance of its monitoring and reporting processes may be based on a view template that is configured and/or customized by the user. In one embodiment, such a view may be considered a dashboard, and its template a dashboard template. The following description relates to embodiments of a user interface whereby a user may create, configure, and/or customize such a view template for asset groups.

[0344] FIG. 49 illustrates a user interface display for adding an asset group view for an asset group. User interface display **6000** is such as might be caused to display during the processing of block **4020** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. User interface display **6000** is the same as interface **4700** of FIG. 36, with the addition of add asset group view window **6002** displayed in response to user selection of the add view action element **4706** in interface **4700**. Add asset group view window **6002** includes a group view name box **6004** and a

group view description box **6004**. Boxes **6002** and **6004** are initially presented to the user as blank input boxes. Boxes **6002** and **6004** are editable to allow a user to input information regarding the new group view in boxes **6002**, **6004**. As illustrated in interface **6000** a new asset group view title “Fuel Consumption” is provided. Selection of the cancel button **6010** causes the add group view window **6002** to be closed without any changes saved for the new asset group view. Selection of the add button **6020** causes the new group view to be added to the asset group with the information indicated in the group view name box **6004** and group view description box **6006**.

[0345] FIG. 50 illustrates a user interface display for configuring an asset group view via an asset group view console. User interface display **6100** is such as might be caused to display during the processing of block **4020** of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as **2134** of FIG. 10. In one embodiment, interface **6100** is displayed in response to selection of the add button **6020** of interface **6000** of FIG. 49. Interface **6100** of FIG. 50 illustrates one possible embodiment of a user interface whereby a user may create, configure, and/or customize a view template for an asset group.

[0346] User interface display **6100** of FIG. 50 is shown to include system header bar **2402**, application information and menu bar **6102**, toolbar **6140**, display mode action button **6142**, group hierarchy display area **6104**, group view template display area **6106**, and group view configuration area **6108**. System header bar **2402** is as described for identically numbered elements appearing in, and described in relation to, depictions of user interface displays in earlier figures. Application information and menu bar **6102** is comparable to other application information and menu bars depicted and described in relation to earlier appearing figures, such as bar **2602** of FIG. 17. Application information and menu bar **6102** may include an identification of the group view (e.g., “Fuel Consumption”) being configured in interface **6100**.

[0347] Application information and menu bar **6102** may also include various action buttons including, but not limited to, a timeframe action button **6120**, a Preview action button **6122** a Clear action button **6123**, a Revert action button **6124**, a Close action button **6125**, and a Save action button **6126**. These various action buttons may be used in an interface associated with the function of creating, configuring, and/or customizing a view template for an asset group. The action button interface components of this illustrative example are interactive elements that enable a user to make indications of data values and desired actions, for example, to AMRS functionality, which the AMRS computing machinery can then process according to its design, circuitry, and programming.

[0348] Timeframe action button **6120** may be a drop down selection list interface component that enables a user to indicate a desired time or time frame of data to use in relation to data-driven or data-aware elements that may be included in a view template. A user interaction with Preview button **6122**, when selected, displays a preview of the configured view for the asset group. A user interaction with Clear action button **6123** in an embodiment may cause the view template under construction to be emptied of all of its content. A user interaction with Revert action button **6124** in an embodiment may have the effect of causing recent changes made to a view template to be abandoned. In one

embodiment a user interaction with Revert action button **6124** may cause the abandonment of only the single most recent change to the view template. In one embodiment, a user interaction with Revert action button **6124** may cause the abandonment of all changes made to the view template since the last time Save action button **6126** was activated. In one embodiment, a user interaction with Revert action button **6124** may cause the abandonment of all changes made to the view template since the last time an autosave action was performed by the AMRS console processor. Other embodiments are possible. A user interaction with Close action button **6125** may cause the abandonment of all changes made to the view template and to dismissal of the view template (e.g., interface **6100**) altogether. A user interaction with Save action button **6126**, in one embodiment, may cause the current configuration of the view template under construction to be reflected in computer storage anyway such that it may be recalled or restored, perhaps by reflecting the configuration information in a named file in the filesystem of a host computer.

[0349] Group hierarchy display area **6104** is shown to include assets tab **6130**, groups tab **6132**, other tab **6134**, search component **6136**, and asset group list **6138**. Assets tab **6130**, groups tab **6132**, and other tab **6134** has similar content, organization, appearance, and formatting to asset tabs, group tabs, and other tabs of user interfaces illustrated and discussed in relation to earlier appearing figures. Similarly, a working understanding of search component **6136** may be developed by consideration of what has come before. Asset group list **6138** is shown to include asset group list entry representing the asset group named “siemens engine” having asset group metrics named “fuel_consumption” **6160**, “fuel_consumption_AVG” **6170**, “oil_level”, “rpm”, and “temperature”.

[0350] Display mode action button **6142** enables the user to request a toggling action between two alternate display modes, “Grid On” and “Grid Off.” A display mode of “Grid Off” is represented in interface **6100** of FIG. 50 as it appears and provides user interface components for creating, configuring, and/or customizing (i.e., editing) a view template without an underlying grid format displayed. A user interaction with display mode action button **6142** may result in the transition to a “Grid On” mode user interface display where an underlying grid format is displayed in the group view template display area **6106**.

[0351] Toolbar **6140** is shown as having a number of tool icons such as “T” icon. Icons in toolbar **6140** may be selected by a user interaction to cause a particular effect, engage a particular function, and/or begin a particular operational mode. For example, in one embodiment, a mouse click on the “T” icon may cause the addition of an empty, default-sized text display element to the view template under construction, causing it to appear at a default location in group view template display area **6106**, and engaging an operational mode for entering text into the newly introduced text display element. In an embodiment, many of the tools represented in toolbar **6140** may be associated with adding different types of elements to the group view template, and with manipulating the elements that are present in the template. In one embodiment, one or more tool icons may be associated with static elements that may be included in the design of the group view template. Such static elements may include text blocks or labels, imported graphical imagery

(e.g., icons, picture files, videos, fixed animations), or drawing elements such as shapes and lines.

[0352] In one embodiment, one or more tool icons may be associated with dynamic elements that may be included in the design of the group view template. Such dynamic elements may be data-driven or data-aware and may determine one or more aspects of their appearance or behavior at a point in time based on currently supplied data. Such data-driven or data-aware dynamic elements may be referred to as “widgets” in one embodiment. In one AMRS embodiment, a monitoring/reporting processor, and CCC console processor functions related thereto, may include functionality to implement a number of built-in widgets and may further include functionality to implement an extensible widget framework which functionality may include, for example, functionality to recognize, install, or activate widget modules, and functionality to exercise the content of those modules. In one such embodiment, widget modules may be packaged after the fashion of programming objects and have attributes or properties (associated data) and methods or behaviors (programmed actions) which may be accessible and/or exercisable by a recognized interface. In one embodiment, a supported widget may be limited to receiving a single data factor or component that drives it, such as the data of a particular metric for a particular asset. In one embodiment, a supported widget may be able to receive multiple data factors or components to drive it, such as the data of different metrics that may be associated with the same asset. These and other embodiments are possible.

[0353] Group view template display area 6106 is shown to include a representation of some or all of the group view template currently under construction. (For example, a representation of only some of the current group 6106 view template may appear in display area 6106 where display area 6016 is smaller than the size of the current group view template. In such a case, group view template display area 6106 may be scrollable.) The current group view template is shown to include static graphical elements including, for example, a diesel engine icon/shape/picture 6155. The current group view template is further shown to include multiple dynamic elements including widgets 6165 and 6175. Widget 6165 is used to depict a current fuel consumption asset group metric. Widget 6175 is used to depict an average (AVG) fuel consumption asset group metric. Data related to the “fuel_consumption” group metric 6160 and “fuel_consumption_AVG” group metric 6170 may drive the appearance of the widget 6165, 6175, and, particularly, the current value for each of the metrics is displayed in a corresponding one of the metric presentation blocks or tiles of the widget, and an urgency level associated with each of the values determines the color of the text used to display the current value in the metric presentation blocker tile (for example, the value of 2.47 for the fuel consumption metric 6160 may be associated with a normal level and so may display in green).

[0354] Configuration display area 3270 may include (not shown), but are not limited to, display alerts area where defined alerts (e.g., as defined during processing of block 4018 of FIG. 29) to be displayed in association with the widget are represented by tokens; general drawing attributes area; title display option area; data-driven animation control section; drilldown area where a user interface navigation target, such as a Diagnose interface display, is defined to use in circumstances where a user double clicks or performs

some other specified interaction with the displayed widget; and update action button which enables a user to indicate the desire to synchronize the representation of the widget displayed in with the representation of its configuration information displayed.

[0355] It is noted that a widget-depicted group metric may have a corresponding asset group metric entry visible in group hierarchy display area 6104, or not—as in the case where an asset group is in a collapsed state in asset group list 6138. For example, the fuel consumption widget 6165 may correspond to fuel_consumption group metric 6160 and the AVG consumption widget 6175 may correspond to the fuel_consumption_AVG group metric 6170 in the asset group list 6138.

[0356] Asset selection action element 6150 may be a drop down selection list interface component that enables a user to indicate a desired asset node that is a member of the asset group indicated as active in the group hierarchy display area 6014 that may be included in the group view template display area 6106. Asset selection action element 6150 may include a list of all asset node members of the asset group to select from. When the creation, construction, editing, or such for a view template is complete, a user may activate the Save action button 6126 to safely fix a definition/configuration of the group view template in asset group view browse interface 6200 described with respect to FIG. 51A.

[0357] FIGS. 51A-51C illustrate user interface displays for asset group view browsing at the asset group level. Referring to FIG. 51A, user interface display 6200 is such as might be caused to display during the processing of block 4020 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 6200 is the same as interface 4700 of FIG. 36, with the addition of a view tab 6202 for the fuel consumption (referred to a Fuel Consumption view tab 6202) displayed in response to user selection of the Save action button 6126 of FIG. 50.

[0358] Interface 6200 illustrates that Fuel Consumption view tab 6202 as the active tab. The Fuel Consumption view tab 6202, when activated, displays the group view configured via interface 6100 of FIG. 50. The group view displayed in interface 6200 may be controlled and/or changed via asset selection action element 6150. As previously discussed, asset selection action element 6150 may be a drop down selection list interface component that enables a user to indicate a desired asset node that is a member of the asset group indicated as active in the group view of interface 6200. Asset selection action element 6150 may include a list of all asset node members of the asset group to select from. As shown in interface 6200 of FIG. 51A, asset node “series 100.100_1100.s_engine” 6205 is currently selected as the active asset in the group view of interface 6200.

[0359] The current group view of interface 6200 is further shown to include multiple dynamic elements including widgets 6165 and 6175. As discussed above, widget 6165 is used to depict a current fuel consumption asset group metric. Widget 6175 is used to depict an average (AVG) fuel consumption asset group metric. Data related to the fuel_consumption group metric and fuel_consumption_AVG group metric may drive the appearance of the widget 6165, 6175 in interface 6200, and, particularly, the current value for each of the metrics is displayed in a corresponding one of the metric presentation blocks or tiles of the widget, and an urgency level associated with each of the values deter-

mines the color of the text used to display the current value in the metric presentation blocker tile (for example, the value of 2.47 for the fuel consumption metric may be associated with a normal level and so may display in green). [0360] FIG. 51B illustrates interface 6210, which is the same as interface 6200 of FIG. 51A, with the distinction of the asset selection action element 6150 being set to select asset “series_100.100_1101.s_engine” 6215 of the asset group “siemens engine.” In an embodiment, asset 6215 may be displayed in response to user interaction with the asset selection action element 6150. The current group view of interface 6210 is shown to include widgets 6165 and 6175. As discussed above, the current value for each of the corresponding metrics of the selected asset 6215 of the asset group is displayed in the widgets 6165, 6175, where with an urgency level associated with each of the values determines the color of the text used to display the current value in the metric presentation blocker tile (for example, the value of 8.35 for the fuel consumption metric may be associated with a low urgency or caution level and so may display in orange).

[0361] FIG. 51C illustrates interface 6220, which is the same as interface 6200 of FIG. 51A ad interface 6210 of FIG. 51B, with the distinction of the asset selection action element 6150 being set to select asset “series_100.100_1102.s_engine” 6225 of the asset group “siemens engine.” In an embodiment, asset 6225 may be displayed in response to user interaction with the asset selection action element 6150. The current group view of interface 6210 is shown to include widgets 6165 and 6175. As discussed above, the current value for each of the corresponding metrics of the selected asset 6225 of the asset group is displayed in the widgets 6165, 6175, where with an urgency level associated with each of the values determines the color of the text used to display the current value in the metric presentation blocker tile (for example, the value of 10.13 for the fuel consumption metric may be associated with a high urgency level and so may display in red).

[0362] FIGS. 52A-52C illustrate user interface displays for asset group view browsing at the asset level. User interface display 6300, 6310, and 6320 are such as might be caused to display during the processing of block 4020 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interfaces 6300, 6310, 6320 are the same as respective interfaces 6200, 6210, and 6220 of FIGS. 51A-51C. However, interface 6300, 6310, and 6320 are asset group view browsing interfaces from the perspective of the asset tab 4222 (e.g., see interface 4200 of FIG. 31).

[0363] With respect to interface 6300 of FIG. 52A, when asset node “100_series.100_1100.s_engine” 6302 is selected as the active asset node under the asset tab 4222, the corresponding group view for the asset node 6302 is displayed under the view tab 6304. Interface 6300 illustrates that Fuel Consumption view tab 6302, when activated, displays the group view configured via interface 6100 of FIG. 50. The current group view of interface 6300 is further shown to include multiple dynamic elements including widgets 6165 and 6175, which is the same as discussed above with respect to interface 6200 of FIG. 51A.

[0364] With respect to interface 6310 of FIG. 52B, when asset node “100_series.100_1101.s_engine” 6312 is selected as the active asset node under the asset tab 4222, the corresponding group view for the asset node 6312 is dis-

played under the view tab 6304. Interface 6300 illustrates that Fuel Consumption view tab 6302, when activated, displays the group view configured via interface 6100 of FIG. 50. The current group view of interface 6300 is further shown to include multiple dynamic elements including widgets 6165 and 6175, which is the same as discussed above with respect to interface 6210 of FIG. 51B.

[0365] With respect to interface 6320 of FIG. 52C, when asset node “100_series.100_1102.s_engine” 6322 is selected as the active asset node under the asset tab 4222, the corresponding group view for the asset node 6322 is displayed under the view tab 6304. Interface 6300 illustrates that Fuel Consumption view tab 6302, when activated, displays the group view configured via interface 6100 of FIG. 50. The current group view of interface 6300 is further shown to include multiple dynamic elements including widgets 6165 and 6175, which is the same as discussed above with respect to interface 6220 of FIG. 51C.

[0366] FIG. 53 illustrates a user interface display for asset group view monitoring configuration console. User interface display 6400 is such as might be caused to display during the processing of block 4020 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 6400 may cause to be displayed when monitor selector 4212 (see interface 4200 of FIG. 31) is selected as the active operational selector.

[0367] Interface 6400 provides a monitoring view overview. Interface 6400 includes a view list 6420 of all views configured for asset groups. The view list 6420 may be in a tabular format and can include identifying information for each asset group. The identifying information may include, but is not limited to, a name 6402, a group 6404, an asset 6406, and actions 6408. Interface 6400 may also include a search component 6410 to allow for filtering of asset groups in the view list 6420. A working understanding of search component 6410 may be developed by consideration of what has come before.

[0368] As illustrated in interface 6400, view list 6420 is shown with a group view entry having a name 6402 of “Fuel Consumption”, which was configured via interface 6100 of FIG. 50. Group 6404 identifies group view “Fuel Consumption” as part of the “Siemens engine” asset group. No assets are displayed in the asset 6406 field. Actions 6408 for the “Fuel Consumption” group view may include, but are not limited to, Rename and Delete. The items listed in action 6408 field may be actionable elements that, when selected by a user, cause processing to occur effect the indicated action. The Add action button 6412, when selected by a user, may cause a view configuration interface, such as interface 6100 of FIG. 50, to be displayed to allow a new group view to be configured and defined. In other embodiments, the Add action button 6412, when selected by a user, may cause a new view configuration interface for an individual asset view to be displayed.

[0369] FIGS. 54A-54C illustrate user interface displays for asset group view monitoring. User interface display 6500, 6510, and 6520 are such as might be caused to display during the processing of block 4020 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interfaces 6500, 6510, 6520 are the same as respective interfaces 6200, 6210, and 6220 of FIGS. 51A-51C. However, interface 6500, 6510, and 6520 are asset group view

monitoring interfaces that are displayed when the monitor selector 4212 is selected as the active operational selector. [0370] With respect to interface 6500 of FIG. 54A, when asset node “100_series.100_1100.s_engine” 6502 is selected as the active asset node 6504 via the asset selection action element 6150, the corresponding group view for the asset node 6502 is displayed. Interface 6500 displays the group view configured via interface 6100 of FIG. 50. The current group view of interface 6500 is further shown to include multiple dynamic elements including widgets 6165 and 6175, which is the same as discussed above with respect to interface 6200 of FIG. 51A.

[0371] With respect to interface 6510 of FIG. 54B, when asset node “100_series.100_1101.s_engine” 6512 is selected as the active asset node 6514 via the asset selection action element 6150, the corresponding group view for the asset node 6512 is displayed. Interface 6500 displays the group view configured via interface 6100 of FIG. 50. The current group view of interface 6500 is further shown to include multiple dynamic elements including widgets 6165 and 6175, which is the same as discussed above with respect to interface 6200 of FIG. 51A.

[0372] With respect to interface 6520 of FIG. 54C, when asset node “100_series.100_1102.s_engine” 6522 is selected as the active asset node 6524 via the asset selection action element 6150, the corresponding group view for the asset node 6522 is displayed. Interface 6500 displays the group view configured via interface 6100 of FIG. 50. The current group view of interface 6500 is further shown to include multiple dynamic elements including widgets 6165 and 6175, which is the same as discussed above with respect to interface 6200 of FIG. 51A.

[0373] The user interfaces already discussed in relation to group metric monitoring and reporting for an asset group, namely, the user interfaces illustrated and discussed in relation to FIGS. 49-54C, have largely related to interfaces employed by a command, control, and configuration console processor (such as CCC console 2134 of FIG. 10). These interfaces enable a user, such as a system administrator or operator, to manipulate the virtual levers, buttons, dials, and switches (embodied in the information of a CCC data store such as 2132 of FIG. 10) that control the operation of the asset monitoring and reporting machine. The focus now turns with the discussion of the figures that follows to asset group management performed by the AMRS machinery. Figures that follow relate largely to interfaces that may be utilized by the AMRS during the processing described and contemplated for block 4050 of FIG. 29 as performed by a monitor/reporter processor such as 2142 of FIG. 10, for example.

3.4.4 Asset Group Management

[0374] In some embodiments, asset node members of an asset group may not all have the same attributes (e.g., metrics). An advanced configuration and setting interface may be provided that allows a user to configure and manage the asset group when encountering such a situation. The advanced configuration and setting interface may be utilized to configure and manage the asset group in various other use cases as may be encountered during asset group managed performed by the AMRS.

[0375] FIG. 55 illustrates a user interface display for an asset group management console. User interface display 6600 is such as might be caused to display during the

processing of block 4022 of FIG. 29, for example; such processing possibly performed by a command, control, and configuration console processor, such as 2134 of FIG. 10. Interface 6600 is similar to interface 4700 of FIG. 36 and interface 5100 of FIG. 40. Interface 6600 illustrates a group overview tab 6602 for an asset group (e.g., “engine”). Group overview tab 6602 details metrics configured for an asset group (“engine”) that is selected as an active asset group in interface 6600.

[0376] Interface 6600 illustrates that group overview tab 6602 as the active tab. The group overview tab 6602, when activated, provides a group metrics overview display portion of interface 6600. The group metrics overview display includes an asset member list (also referred to as members list) 6604 of the asset group and a metrics list 6606 of the asset group. Interface 6600 indicates that there are 5 assets in the assets list 6604 associated with the “engine” asset group. Interface 6600 further indicates that there are 5 different metrics in the metrics list 6606 associated with the “engine” asset group, where the metrics are “fuel_consumption”, “oil_level”, “rpm”, “tmp”, and “temperature”. Metrics list 6606 may be provided in a tabular format with multiple column headers providing identifying information corresponding to each metric entry in the metrics list 6606. The column headers may include, but are not limited to, metric name 6610, occurrence 6612, assets 6614, share metric 6616, and action 6618. Metric name 6610 indicates the name of the corresponding group metric. Occurrence 6612 provides a relative occurrence of the particular group metric in terms of the total number of asset nodes in the group. For example, the occurrence value of group metric “fuel_consumption” is 100% as 123 out of 123 asset node members of the asset group are configured with the particular metric. On the other hand, the group metric “temperature” has an occurrence value of 3%, as 3 out of 123 asset node members of the asset group are configured with this particular metric. Occurrence 6612 is illustrated with a bar graph to visually represent the value of occurrence. However, other visual aids may be utilized for the occurrence value. Asset 6614 indicates the count of asset members configured with the particular group metric as compared to the total number of asset group members. Share metric 6616 is a toggle action element that, in response to activation by a user, causes the particular group metric to be replicated across of member nodes of the asset group. Action 6618 may include one or more selectable actions that can be performed with respect to a group metric for advanced configuration purposes (e.g., Group, Ungroup, Join, Disjoin, Generate, Analyze, Delete, etc.).

[0377] Actions 6618 of metrics list 6606 may each be enabled for user interaction so as to enable a user to indicate a selection of one of the actions of the list, in an embodiment. The selection of an action 6618 of metrics list 6606 may result in the display of information and user interface elements in another user interface. In one example, the group action (which may also be referred to as a Join action) creates a joint group metric that combines disparately named metrics into a single group metric.

[0378] FIG. 56 illustrates a user interface display for joining group metrics of an asset group. User interface display 6700 is the same as interface 6600 of FIG. 55, with the addition of join group metric window 6702 displayed in response to user selection of a “Group” or “Join” action 6618 from the metrics list 6606 of interface 6600. Join group

metric window 6702 includes a drop-down menu 6704 providing a list of candidate group metrics that are joinable with the group metric for which the “Group” action 6618 was selected. As illustrated in interface 6700, the group metric “temperature” is selected for joining with the group metric “tmp”. Selection of the Group action button 6706 causes a joint group metric to be generated and saved.

[0379] FIG. 57 illustrates a user interface display for joint group metrics of an asset group. User interface display 6800 is the same as interface 6700 of FIG. 55, with the addition of new joint group metric “temperature 6802 displayed in response to user selection of the “Group” action button 6706 in interface 6700 of FIG. 56. Joint group metric temperature 6802 is illustrated as an expandable element that when expanded displays the constituent group metrics of the joint group metric. The constituent group metrics of the joint group metric “temperature” include the “tmp” group metric 6804 and the “temperature” group metric 6806. The Actions column 6618 for a joint group metric can also include the “Ungroup” or “Disjoin” action 6808 that, when selected by user, causes the corresponding group metric to be removed from the joint group metric.

[0380] The user interfaces already discussed in relation to asset group management, namely, the user interfaces illustrated and discussed in relation to FIGS. 55-57, have largely related to interfaces employed by a command, control, and configuration console processor (such as CCC console 2134 of FIG. 10). These interfaces enable a user, such as a system administrator or operator, to manipulate the virtual levers, buttons, dials, and switches (embodied in the information of a CCC data store such as 2132 of FIG. 10) that control the operation of the asset monitoring and reporting machine.

[0381] The preceding description sets forth numerous specific details such as examples of specific systems, components, methods, and so forth, in order to provide a good understanding of several embodiments of the present invention. It will be apparent to one skilled in the art, however, that at least some embodiments of the present invention may be practiced without these specific details. In other instances, well-known components or methods are not described in detail or are presented in simple block diagram format in order to avoid unnecessarily obscuring the present invention. Thus, the specific details set forth are merely exemplary. Particular implementations may vary from these exemplary details and still be contemplated to be within the scope of the present invention.

[0382] In the above description, numerous details are set forth. It will be apparent, however, to one of ordinary skill in the art having the benefit of this disclosure, that embodiments of the invention may be practiced without these specific details. In some instances, well-known structures and devices are shown in block diagram form, rather than in detail, in order to avoid obscuring the description.

[0383] Some portions of the detailed description are presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory. These algorithmic descriptions and representations are the means used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of steps leading to a desired result. The steps are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or

magnetic signals capable of being stored, transferred, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

[0384] It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. Unless specifically stated otherwise as apparent from the above discussion, it is appreciated that throughout the description, discussions utilizing terms such as “determining”, “identifying”, “adding”, “selecting” or the like, refer to the actions and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (e.g., electronic) quantities within the computer system’s registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage, transmission or display devices.

[0385] Embodiments of the invention also relate to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program may be stored in a computer readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROM s, and magnetic-optical disks, read-only memories (ROM s), random access memories (RAM s), EPROMS, EEPROM s, magnetic or optical cards, or any type of media suitable for storing electronic instructions.

[0386] The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general purpose systems may be used with programs in accordance with the teachings herein, or it may prove convenient to construct a more specialized apparatus to perform the required method steps. The required structure for a variety of these systems will appear from the description below. In addition, the present invention is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages may be used to implement the teachings of the invention as described herein.

[0387] Implementations that are described may include graphical user interfaces (GUIs). Frequently, an element that appears in a GUI display is associated or bound to particular data in the underlying computer system. The GUI element may be used to indicate the particular data by displaying the data in some fashion, and may possibly enable the user to interact to indicate the data in a desired, changed form or value. In such cases, where a GUI element is associated or bound to particular data, it is a common shorthand to refer to the data indications of the GUI element as the GUI element, itself, and vice versa. The reader is reminded of such shorthand and that the context renders the intended meaning clear to one of skill in the art where a distinction between a GUI element and the data to which it is bound is meaningful.

[0388] It is to be understood that the above description is intended to be illustrative, and not restrictive. Many other embodiments will be apparent to those of skill in the art upon reading and understanding the above description. The scope of the invention should, therefore, be determined with

reference to the appended claims, along with the full scope of equivalents to which such claims are entitled.

[0389] The preceding point may be elaborated with a few examples. Many details have been discussed and disclosed in regards to user interfaces including graphical user interfaces (GUIs). While it is convenient to describe inventive subject matter in terms of embodiments that include familiar technologies, components, and elements, the inventive subject matter should not be considered to be constrained to these, and the ready availability and appropriateness of substitutes, alternatives, extensions, and the like is to be recognized. What may be shown or described as a single GUI or interface component should liberally be understood to embrace combinations, groupings, collections, substitutions, and subdivisions in an embodiment. What may be shown or described as a single GUI or interface component may be embodied as an atomic or truly elemental interface component, or may readily be embodied as a complex or compound component or element having multiple constituent parts. What may be shown, described, or suggested to be a uniformly shaped and contiguous GUI or interface component, such as an interface region, area, space, or the like, may be readily subject to implementation with non-uniformly shaped or noncontiguous display real estate.

[0390] As yet one more example, apparatus that perform methods, processes, procedures, operations, or the like, disclosed herein may be referred to as a computer, computer system, computing machine, or the like. Any such terminology used herein should be reasonably understood as embracing any collection of temporarily or permanently connected hardware devices in combination with any software each requires to operate and perform operations and functions necessary to an implementation of an inventive aspect. Adopting such an understanding is consistent with modern computing practices and eliminates the need to obscure the disclosure of inventive aspects with catalogs of implementation options and alternatives.

[0391] As one final example, methods, procedures, or processes may be described herein by reference to flow charts or block diagrams and possibly in terms of sequences of steps or operations. It should be understood, however, that the practice of an inventive aspect is generally not limited to the number, ordering, or combination of operations as may be described for an illustrative embodiment used to teach and convey an understanding of inventive aspects possibly within a broader context. Accordingly, not all operations or steps described are illustrated may be required to practice of an inventive aspect. Different embodiments may variously omit, augment, combine, separate, reorder, or reorganize the performance of operations, steps, methods, procedures, functions, and the like disclosed or suggested herein without departing from an inventive aspect. Further, where sequences of operations may be illustrated, suggested, expressed, or implied, an embodiment practicing inventive aspects may perform one or more of those operations or sets of operations in parallel rather than sequentially.

[0392] Accordingly, inventive aspects disclosed herein should be considered broadly without unnecessary limitation by the details of the disclosure, and should be considered as limited only by accompanying claims or where reason demands it.

What is claimed is:

1. A method performed by one or more processing devices, the method comprising:

causing display of a user interface for managing an asset group, wherein the user interface comprises:

a set of assets associated with the asset group, the asset group associating a plurality of assets of an asset hierarchy as a group, and

a set of group metrics associated with the asset group, wherein each group metric of the set of group metrics is derived from a search of machine data associated with the asset group, and wherein at least a portion of the assets associated with the asset group correspond with different group metrics of the set of group metrics.

2. The method of claim 1, wherein the user interface for managing the asset group is displayed based on a selection of an indication of the asset group.

3. The method of claim 1, wherein the set of group metrics are displayed in a tabular format with a plurality of columns providing identifying information corresponding to each group metric of the set of group metrics.

4. The method of claim 1, wherein each group metric of the set of group metrics is displayed in association with an occurrence of a corresponding group metric relative to a total number of assets in the set of assets.

5. The method of claim 1, wherein each group metric of the set of group metrics is displayed in association with a proportion of assets of the set of assets configurated with the corresponding group metric.

6. The method of claim 1, wherein each group metric of the set of group metrics is displayed in association with a visual representation indicating a value of assets of the set of assets configured with the corresponding group metric.

7. The method of claim 1, wherein each group metric of the set of group metrics is displayed in association with a count of assets of the set of assets configured with the corresponding group metric.

8. The method of claim 1, wherein each group metric of the set of group metrics is displayed in association with a toggle action element that, upon selection, causes the corresponding group metric to be replicated across at least one other asset of the set of assets associated with the asset group.

9. The method of claim 1, wherein each group metric of the set of group metrics is displayed in association with one or more selectable actions for an advanced configuration for the corresponding group metric.

10. The method of claim 1, wherein at least one group metric of the set of group metrics is displayed in association with a selectable action that, upon selection, causes performance of a group action, an ungroup action, a join action, a disjoin action, a generate action, an analyze action, a delete action, or a combination thereof.

11. The method of claim 1, wherein a group metric of the set of group metrics is displayed in association with a selectable action, and wherein when the selectable action is selected, a new user interface is presented including information association with the selectable action.

12. The method of claim 1, wherein a group metric of the set of group metrics is displayed in association with a selectable action comprising a join action, and wherein when the join action is selected, a joint group metric is created that combines disparate group metrics into a single group metric.

- 13.** A system comprising:
a memory; and
one or more processing devices coupled with the memory,
the one or more processing devices configured to:
cause display of a user interface for managing an asset group, wherein the user interface comprises:
a set of assets associated with the asset group, the asset group associating a plurality of assets of an asset hierarchy as a group, and
a set of group metrics associated with the asset group, wherein each group metric of the set of group metrics is derived from a search of machine data associated with the asset group, and wherein at least a portion of the assets associated with the asset group correspond with different group metrics of the set of group metrics.
- 14.** The system of claim **13**, wherein the user interface for managing the asset group is displayed based on a selection of an indication of the asset group.
- 15.** The system of claim **13**, wherein the set of group metrics are displayed in a tabular format with a plurality of columns providing identifying information corresponding to each group metric of the set of group metrics.
- 16.** The system of claim **13**, wherein each group metric of the set of group metrics is displayed in association with an occurrence of a corresponding group metric relative to a total number of assets in the set of assets.
- 17.** A non-transitory computer-readable storage medium encoding executable instructions thereon that, in response to execution by one or more processing devices, cause the one or more processing devices to:
- cause display of a user interface for managing an asset group, wherein the user interface comprises:
a set of assets associated with the asset group, the asset group associating a plurality of assets of an asset hierarchy as a group, and
a set of group metrics associated with the asset group, wherein each group metric of the set of group metrics is derived from a search of machine data associated with the asset group, and wherein at least a portion of the assets associated with the asset group correspond with different group metrics of the set of group metrics.
- 18.** The non-transitory computer-readable storage medium of claim **17**, wherein each group metric of the set of group metrics is displayed in association with one or more selectable actions for an advanced configuration for the corresponding group metric.
- 19.** The non-transitory computer-readable storage medium of claim **17**, wherein at least one group metric of the set of group metrics is displayed in association with a selectable action that, upon selection, causes performance of a group action, an ungroup action, a join action, a disjoin action, a generate action, an analyze action, a delete action, or a combination thereof.
- 20.** The non-transitory computer-readable storage medium of claim **17**, wherein a group metric of the set of group metrics is displayed in association with a selectable action comprising a join action, and wherein when the join action is selected, a joint group metric is created that combines disparate group metrics into a single group metric.

* * * * *