

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265015

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Nijssen; Raymond et al.

FPGA MEMORY WITH AUTO ADDRESS MODE

Abstract

Memory blocks often consume many switch box resources. For example, an 8-bit memory using 8-bit addressing uses at least eight address lines, eight data lines, a write-enable line, and a read-enable line. Using an auto address mode, the address and data are multiplexed on the same lines. The initial address and a stride are provided before the writing process begins. Between writes, the address is incremented by the stride. Thus, the memory block is able to determine the address for the next write based on the starting address and the stride, and does not need to receive the new address on the address lines. The auto address mode may be implemented by including a logic block within the memory block. The logic block may be programmed for purposes other than an auto address mode.

Inventors: Nijssen; Raymond (San Jose, CA), Pugh; Daniel (Los Gatos, CA), Akesh; Naveen Bharathwaj (Los Gatos, CA), LaFrieda; Christopher C. (Ridgefield, NJ), Radhakrishnan; Sriram (Santa Clara, CA), Mehta; Ronak (Santa Clara, CA), Ekanayake; Virantha Namal (Baltimore, MD)

Applicant: Achronix Semiconductor Corporation (Santa Clara, CA)

Family ID: 1000007724214

Appl. No.: 18/581131

Filed: February 19, 2024

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

CPC G06F3/0659 (20130101); G06F3/0604 (20130101); G06F3/064 (20130101); G06F3/0679 (20130101);

Background/Summary

TECHNICAL FIELD

[0001] The subject matter disclosed herein generally relates to methods, systems, and machine-readable storage media for providing memory with auto address modes inside a Field Programmable Gate Array (FPGA).

BACKGROUND

[0002] FPGAs comprise logic blocks and memory blocks coupled by switch boxes. Memory blocks often use many wires from switch boxes for address lines, data lines, and control lines. This limits the number of switch box wires that are available for other purposes.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0003] Various of the appended drawings merely illustrate example embodiments of the present disclosure and cannot be considered as limiting its scope.

[0004] FIG. 1 illustrates the architecture of an FPGA with switch boxes, logic blocks, and memory blocks, according to some example embodiments.

[0005] FIG. 2 illustrates a switch box connected to a memory block with auto address modes, according to some example embodiments.

[0006] FIG. 3 illustrates logic to implement an auto address mode in a memory block, according to some example embodiments.

[0007] FIG. 4 illustrates an example logic group suitable for use in implementing the logic of FIG. 3, according to some example embodiments.

[0008] FIG. 5 illustrates a logic block in a memory block, the logic block comprising six of the logic groups of FIG. 4, according to some example embodiments.

[0009] FIG. 6 is a flowchart of a method for implementing an auto address mode in a memory block, according to some example embodiments.

[0010] FIG. 7 is a block diagram illustrating an example of a machine upon or by which one or more example process embodiments described herein may be implemented or controlled.

DETAILED DESCRIPTION

[0011] Example methods, systems, and computer programs are directed to memory blocks in FPGAs that include logic to support auto addressing modes. Examples merely typify possible variations. Unless explicitly stated otherwise, components and functions are optional and may be combined or subdivided, and operations may vary in sequence or be combined or subdivided. In the following description, for purposes of explanation, numerous specific details are set forth to provide a thorough understanding of example embodiments. It will be evident to one skilled in the art, however, that the present subject matter may be practiced without these specific details.

[0012] An FPGA (or other programmable integrated circuit) may include a number of columns. Each column is an array of elements including a switch box and a tile. A column comprising logic block tiles may be referred to as a logic column. A column comprising memory block tiles may be referred to as a memory column.

[0013] The routing of the switch boxes is determined when the FPGA is programmed. After programming, the switch boxes statically connect wires to neighboring switch boxes or the tile. By connecting to other switch boxes, signals can be routed to other tiles or columns. The tile may be a logic block, a memory block, or other specialized logic.

[0014] Memory columns often consume many switch box resources. For example, an 8-bit memory using 8-bit addressing uses at least eight address lines, eight data lines, a write-enable line, and a

read-enable line. This leaves few switch box resources available for other uses, such as connecting to an adjacent column. As a result, longer routes may be used for connections, increasing the latency of the FPGA design.

[0015] Memory is accessed in response to access commands (e.g., reads and writes) that include an address. The addressed portion of memory is accessed. Write commands additionally include a data value to be written to the addressed memory. Write commands are responded to with an acknowledgement that the write was successful. Read commands are responded to with a data value read from the addressed memory. A common use of memory blocks is to perform sequential reads and writes. Sequential accesses are separated by a stride, which refers to the number of bytes it takes to step between accessed elements. For example, one hundred consecutive bytes may be written by performing one hundred consecutive writes in which the address for each write is one more than the previous write, using a stride of one.

[0016] Using an auto address mode, the address and data are multiplexed on the same lines. The initial address and a stride are provided before the writing process begins. Between writes, the address is incremented by the stride. Thus, the memory block is able to determine the address for the next write based on the starting address and the stride and does not need to receive the new address on the address lines. As a result, fewer lines are used from the switch box to the memory block, allowing more lines to be used for other purposes and, depending on the particular design, reducing overall latency.

[0017] The auto address mode may be implemented by including a logic block within the memory block. The logic block may be programmed for purposes other than an auto address mode. For example, a design that does not make use of all memory blocks may use the logic block within the unused memory blocks to implement additional application logic. This increases the overall logic density of the FPGA. Alternatively, the logic block may implement additional functionality related to the memory block. For example, the logic block may remap addresses dynamically in the event of faulty portions of the memory. As another example, the logic block can compare addresses on the read and write ports to implement a first-in, first-out (FIFO) queue with almost empty and almost full flags.

[0018] FIG. 1 illustrates the architecture of an FPGA **100** with switch boxes, logic blocks, and memory blocks, according to some example embodiments. The FPGA **100** includes logic columns **120A** and **120B** and memory columns **130A** and **130B**. The logic columns **120A-120B** may be referred to collectively as logic columns **120** or generically as a logic column **120**. The memory columns **130A-130B** may be referred to collectively as memory columns **130** or generically as a memory column **130**. The logic columns **120** comprise switch boxes and logic blocks. The memory columns **130** comprise switch boxes and memory blocks. Switch boxes **140A** and **140B**, logic blocks **150A** and **150B**, and a memory block **160** will be referred to as examples of each type of component.

[0019] Each switch box comprises multiple horizontal and vertical channels. The switch box is configured to connect or disconnect the channels where they cross. Thus, the switch box **140A** can be configured to connect wires from the west side of the FPGA **100** to the logic block **150A** (using a convention that the top of FIG. 1 is the north direction). Alternatively, the switch box **140A** can be configured to connect one or more wires from the west side of the FPGA **100** to vertical wires connecting the switch box **140A** with the switch box **140B**. The switch box **140B** can be configured to connect those vertical wires to the logic block **150B** or to allow the signals on those wires to continue propagating north. Thus, the column of switch boxes in the logic column **120A** allows for the interconnection of any incoming wire from the west of the FPGA **100** to any logic block in the logic column **120A**.

[0020] Similarly, the column of switch boxes in the memory column **130A** allows the data signals from the logic blocks of the logic column **120A** to be routed to any of the memory blocks of the memory column **130A**, such as the memory block **160**. The inputs to the memory blocks may

include address lines, data lines, a data-ready signal, a write indicator, a read indicator, or any suitable combination thereof. In some example embodiments, channels skip some of the switch boxes. For example, a vertical channel may pass through every other switch box in a column. As another example, a hex-based design may be used instead of a square-based design and the hexes may connect to every sixth switch box.

[0021] As discussed herein, the memory block **160** may include an integrated logic block that supports one or more auto address modes. When using an auto address mode, the number of wires dedicated to the memory block **160** is reduced when compared to the normal address mode. In the normal address mode, an address is provided to the memory block **160** each time data is accessed. In the auto address mode, an address and a stride are loaded into the memory block **160**, and further commands may use a 1-bit signal to indicate that the address should be increased by the stride for the command. Data lines can be used instead of address lines during loading of the address and stride. Accordingly, the routing resources of the switch boxes in the memory columns **130** are saved for other purposes.

[0022] FIG. 2 illustrates a switch box **210** connected to a memory block **220** with auto address modes, according to some example embodiments. The memory block **220** includes an input crossbar **230**, an output crossbar **240**, a logic block **250**, and memory **260**.

[0023] The switch box **210** provides input signals to the input crossbar **230** and receives output signals from the output crossbar **240**. The input crossbar **230** is configured during the programming of an FPGA to connect specific wires from the switch box **210** to inputs for the logic block **250**, the memory **260**, or both. For example, in a configuration in which auto address mode is not used, the input crossbar **230** may be configured to route address, data, and command signals from the switch box **210** to the memory **260**, bypassing the logic block **250** entirely. As another example, in a configuration in which the memory block **220** is used only to execute logic, the input crossbar **230** may be configured to route command and data signals to the logic block **250**, ignoring the memory **260**. The outputs of the logic block **250**, the memory **260**, or both, are provided to the output crossbar **240** and thereby routed back to the switch box **210**.

[0024] In a configuration supporting an auto address mode, command and data signals may be routed to the logic block **250**, and outputs from the logic block **250** may be routed to the address, data, and command inputs of the memory **260**. The logic block **250** may be implemented using arithmetic logic units (ALUs), look-up tables (LUTs), registers, multiplexers, or any suitable combination thereof. Integration of the logic block **250** into the memory block **220** allows the switch box **210**, the input crossbar **230**, and the output crossbar **240** to be shared between the logic block **250** and the memory **260**. Most of the area overhead of a logic block tile (e.g., the logic block **150A** of FIG. 1) is in the routing. The memory tile already pays the area overhead for the switch box **210** and the crossbars **230-240**. Accordingly, the additional area used by adding the logic block **250** to the memory block **220** is minimal.

[0025] The logic block **250** is designed to include enough resources to implement auto address logic and also has the flexibility to implement additional functionality. For example, the logic block **250** may be configured to remap addresses dynamically in the event of faulty address rows in the memory **260**. As another example, the logic block **250** may implement address comparison functions on the read and write ports to implement a FIFO with almost empty/full flags.

[0026] The memory block **220** may be part of an FPGA and configured to operate in a first mode, a second mode, or a third mode. Configuration of the memory block **220** may occur at the time the FPGA is programmed. In the first mode, the memory block **220** operates normally, receiving an address for each memory access command. In the second mode, the memory block **220** uses the logic block **250** to generate addresses for read accesses. In the third mode, the memory block **220** uses the logic block **250** to generate addresses for write accesses.

[0027] FIG. 3 illustrates logic **300** to implement an auto address mode in a memory block, according to some example embodiments. The logic **300** includes an auto read address mode signal

305; an auto write addressing mode signal **320**; multiplexers **310**, **315**, **330**, **335**, **340**, **345**, **370**, and **375**; registers **325** and **360**; and an adder **350**. The logic block **250** of the memory block **220** of FIG. 2 may include the logic **300**. To illustrate, reference will be made to the switch box **210** and the memory **260** of FIG. 2.

[0028] When the auto read address mode signal **305** is zero, the multiplexer **310** passes the read address received from a connected switch box to the memory **260**. When the auto write addressing mode signal **320** is zero, the multiplexer **315** passes the write address received from a connected switch box to the memory **260**. Thus, when neither auto address mode signal is asserted, the logic **300** does not affect the operation of the memory block **220**.

[0029] To use the auto read addressing mode, the auto read address mode signal **305** is asserted and the auto write addressing mode signal **320** is not asserted. While the write addressing mode signal **320** is not asserted, the multiplexer **375** selects the 0-bit of the read address from the switch box **210** to use as the LDSTRIDE (load stride) signal. While the write addressing mode signal **320** is asserted, the multiplexer **375** selects the 0-bit of the write address from the switch box **210** to use as the LDSTRIDE signal.

[0030] When the LDSTRIDE signal is zero, the multiplexer **370** routes the output of the stride register **360** to the input of the stride register **360**, thereby not affecting the value stored by the stride register **360**. When the LDSTRIDE signal is one, the multiplexer **370** causes the write data received from the switch box **210** to be stored in the stride register **360**. Thus, the 0-bit of the read or write address controls whether the write data from the switch box **210** is stored in the stride register **360**.

[0031] While the write addressing mode signal **320** is not asserted, the multiplexer **345** selects the 1-bit of the read address from the switch box **210** to use as the INC (increment) signal. While the write addressing mode signal **320** is asserted, the multiplexer **345** selects the 1-bit of the write address from the switch box **210** to use as the INC signal.

[0032] When the INC signal is zero, the multiplexer **340** routes the output of the address register **325** toward the input of the address register **325**. When the INC signal is one, the multiplexer **340** instead uses the output of the adder **350**. The adder **350** generates the sum of the current address stored in the address register **325** with the stride stored in the stride register **360**. Thus, the 1-bit of the read or write address controls whether the current address is incremented by the stride.

[0033] Based on the auto write addressing mode signal **320**, the multiplexer **335** selects between the 2-bit of the read address from the switch box **210** and the 2-bit of the write address from the switch box **210**. The selected bit is used as the LDADDR (load address) signal.

[0034] When the LDADDR signal is zero, the multiplexer **330** routes the output of the multiplexer **340** to the address register **325**. Accordingly, when the address is not being loaded, the value stored in the address register **325** is either unchanged or is incremented by the stride, depending on the INC signal. When the LDADDR signal is one, the multiplexer **330** routes the write data from the switch box **210** to the address register **325**. Thus, the 2-bit of the read or write address controls whether the write data from the switch box **210** is stored in the address register **325**.

[0035] If the auto write addressing mode signal **320** is asserted, the multiplexer **315** causes the value stored in the address register **325** to be provided as the write address to the memory **260**. If the auto read address mode signal **305** is asserted, the multiplexer **310** causes the value stored in the address register **325** to be provided as the read address to the memory **260**. Thus, in the auto addressing mode, the full set of write data wires from the switch box **210** are used, but only three address wires are used (to control the LDSTRIDE, INC, and LDADDR signals). The auto write addressing mode signal **320** and the auto read address mode signal **305** may be constant values determined by the logic being implemented by a particular FPGA and thus connected to ground or the operating voltage without using wires of the switch box **210**.

[0036] The WDATA signal from the switch box may be provided to the WDATA input of the memory **260**, bypassing the logic **300**. As a result, the logic **300** controls the address inputs to the

memory **260** without affecting the data inputs. Alternatively, the WDATA signal received from the switch box may be routed to the WDATA input of the memory **260** by the logic **300**, such that the data to be written is received using the same wires that receive the stride for the stride register **360** and the address for the address register **325**.

[0037] FIG. **4** illustrates an example logic group **400**, suitable for use in implementing the logic **300** of FIG. **3**, according to some example embodiments. The logic **300** of FIG. **3** could be implemented using the multiplexers and registers shown, but could also be implemented with a more versatile set of components. Using more versatile components enables the logic block **250** of FIG. **2** to be used to implement other functionality when the auto addressing mode is not being used.

[0038] The logic group **400** includes a LUT block **410**, an ALU **430**, and a register block **440**. The LUT block **410** includes four 6-input LUTs **420A**, **420B**, **420C**, and **420D**. The LUT block **410** receives six inputs for each 6-input LUT, denoted as A[0 . . . 23]. Additionally, the LUT block **410** can be programmed to treat each 6-input LUT as two 5-input LUTs. A 6-input LUT is programmable to generate a 1-bit output for each different 6-bit input. A 5-input LUT is programmable to generate a 1-bit output for each different 5-bit input. The two 5-input LUTs for a 6-input LUT are referred to as the “low” and “high” 5-input LUTs.

[0039] In some example embodiments, the two 5-input LUTs are given the same inputs. In other example embodiments, a subset of the inputs are the same, and a subset of the inputs are different. To support this, additional input may be accepted by the LUT block **410**. For example, the 6-input LUT **420A** may receive inputs A[0 . . . 5] and MLUT[0]. When the 6-input LUT **420A** is configured to operate as a 6-input LUT, inputs A[0 . . . 5] are used. When the 6-input LUT **420A** is configured to operate as two 5-input LUTs, the low 5-input LUT may operate on A[0 . . . 4], and the high 5-input LUT may operate on A[0 . . . 2], MLUT[0], A[5]. Thus, three inputs to the two 5-input LUTs are the same and two can differ.

[0040] The output of the 6-input LUTs **420A-420D** is shown as L6[0 . . . 3]. The outputs of the eight 5-input LUTs are shown as L5H[0 . . . 3] and L5L[0 . . . 3]. For each one of the 6-input LUTs **420A-420D**, the single corresponding output in the L6[0 . . . 3] is used or the two corresponding outputs in the L5H[0 . . . 3] and L5L[0 . . . 3], depending on the programming of the LUT block **410**.

[0041] The 8-bit ALU **430** receives two 8-bit inputs, shown as X[0 . . . 7] and Y[0 . . . 7], and a carry-in bit from another ALU. The sum of the 8-bit inputs (and carry in) is output as an 8-bit value, SUMN[0 . . . 7] and a carry-out bit. Using the carry-in and carry-out signals, multiple 8-bit ALUs can be connected together to create larger ALUs.

[0042] The register block **440** includes eight registers **450A**, **450B**, **450C**, **450D**, **450E**, **450F**, **450G**, and **450H**. Each of the registers **450A-450H** may correspond to one bit of the output from the ALU **430**, to one of the 6-input LUTs **420A-420D**, to one of the 5-input component LUTs of one of the 6-input LUTs **420A-420D**, or any suitable combination thereof. Logic for each of the registers **450A-450H** may be configured to store a constant (e.g., 0 or 1), the output of the corresponding 6-input LUT, the output of the corresponding 5-input LUT, or the corresponding output from the ALU **430**. The value of the eight registers **450A-450H** (or a constant, LUT output, or ALU output) is provided as REG_OUT[0 . . . 7].

[0043] FIG. **5** illustrates a 16-bit reusable logic block (RLB16) **500** in a memory block, the RLB16 **500** comprising six of the logic groups **400** of FIG. **4**, according to some example embodiments. Using the RLB16 **500** to form the logic block **250** of FIG. **2**, 3 16-bit ALUs, 24 LUT6s, and 48 bits of register storage are available. The logic shown in FIG. **3** can be implemented using one 16-bit adder (to add the stride to the address), two 16-bit registers (to store the stride and the address), and three 8-bit 2:1 muxes (to handle control signals). Each 8-bit mux is implemented using 8 LUT6s. Accordingly, the logic block **250** also has two 16-bit ALUs and 16 bits of register storage available for other uses.

[0044] The RLB16 **500** is composed of two 8-bit reusable logic blocks (RLB8), each of which includes three of the logic groups **400**. The RLB8 **502** includes logic groups **506**, **508**, and **510**. The RLB8 **504** includes logic groups **512**, **514**, and **516**. The carry-in signals of the ALUs of the logic groups **512**, **514**, and **516** are connected to a constant zero signal. The carry-in signals of the ALUs of the logic groups **506**, **508**, and **510** are connected to the carry-out signals of the logic groups **512**, **514**, and **516**, respectively. The carry-out signals of the ALUs of the logic groups **512**, **514**, and **516** are ignored or used as overflow signals.

[0045] The 8-bit outputs of the RLB8s **502** and **504** may be concatenated and treated as the 16-bit outputs of the RLB **16 500**. Thus, the RLB16 **500** may generate three 16-bit ALU results, three 16-bit register results, and three sets of eight LUT6 outputs. The RLB16 may also be configured to treat the LUT6s as pairs of LUT5s, thus producing three sets of sixteen LUT5 outputs.

[0046] As shown in FIG. **2**, the logic block **250** (and thus the RLB16 **500**) may be coupled to the switch box **210** via the input crossbar **230**. Thus, WDATA, WADDR, RADDR, and control signals may be received from the switch box **210** to control operation of the RLB16 **500**. When the memory **260** of FIG. **2** is not being used, or when the RLB16 **500** is implementing functionality other than or in addition to an auto address mode, additional inputs to the RLB16 **500** may be received from the switch box **210**, such as ALU operands, LUT indices, register values, register access commands, or any suitable combination thereof. Outputs generated by the RLB16 **500** may be provided to the output crossbar **240** of FIG. **2**, for routing to the switch box **210** and, from there, to other tiles of the FPGA **100** of FIG. **1**.

[0047] FIG. **6** is a flowchart of a method **600** for implementing an auto address mode in a memory block, according to some example embodiments. The method **600** includes operations **610**, **620**, **630**, and **640**. By way of example and not limitation, the method **600** is described as being performed by the logic block **250** of the memory block **220** of FIG. **2**, implemented using the RLB16 **500** of FIG. **5** composed of the logic groups **400** of FIG. **4**. The RLB16 **500** may implement the logic shown in FIG. **3**.

[0048] In operation **610**, the memory block **220**, comprising the logic block **250** and the memory **260**, receives an address and a stride. For example, one or more of the LUT6s **420A-420D** may be used to implement the multiplexers **310**, **315**, **330**, **335**, **340**, **345**, **370**, and **375** of FIG. **3**. The address register **325** may be implemented using sixteen-bit storage provided by the logic groups **508** and **514**. The stride register **360** may be implemented using sixteen-bit storage provided by the logic groups **506** and **512**. The address and the stride may be received sequentially. For example, at a first time, the LDADDR signal of FIG. **3** may be asserted and an address received on the WDATA signal. In response, the address is stored in the address register **325**. At a second time, the LDSTRIDE signal of FIG. **3** is asserted, and a stride is received on the WDATA signal. In response, the stride is stored in the stride register **360**.

[0049] The logic block **250**, in operation **620**, accesses the memory **260** using the address. For example, after loading the address and the stride, a first memory access (e.g., a read or a write) is performed using the loaded address. With reference to FIG. **3**, the loopback input to the address register **325** is just the output from the address register **325**, unless the INC signal is asserted. Accordingly, if the INC signal is not asserted before a memory access is performed, and the logic block **250** is configured in AUTO WADDR MODE or AUTO RADDR MODE, the read and write addresses provided to the memory **260** are the address received in operation **610**.

[0050] In response to an increment command, the logic block **250** increments, in operation **630**, the address by the stride. Returning to FIG. **3**, the adder **350** adds the stride stored in the stride register **360** to the current address stored in the address register **325**. When an INC signal is received, the multiplexer **340** selects the output of the adder **350**, which may be implemented using the ALUs of the logic groups **506** and **512** as the input to the address register **325**. As a result, when the INC signal is received, the address stored in the address register **325** is incremented by the stride.

[0051] In operation **640**, the logic block **250** accesses the memory **260** using the incremented

address. As discussed above with respect to the operation **620**, the value stored in the address register **325** is used as the read address and write address while the logic block **250** is configured in AUTO WADDR MODE or AUTO RADDR MODE. Thus, the incremented value of the address is used in the next memory access. Further memory accesses may be performed at addresses separated by the stride by continuing (or repeating) the assertion of INC and performing additional reads or writes.

[0052] Thus, by use of the method **600**, the memory block **220** is enabled to access a sequence of addresses by loading an address and a stride once, then receiving a 1-bit INC signal. The address and the stride may be received using the same wires (e.g., the WDATA wires from the switch box, as shown in FIG. 3). Furthermore, when the memory access is a write, the data to be written may also be received using the same wires as the address and the stride. This reduces the use of interconnections over standard implementations that load the address for each memory access command.

[0053] Example 1 is a memory block comprising: a memory; and a logic block coupled to the memory and configured to perform operations comprising: sequentially receiving an address and a stride using the same wires; accessing the memory using the address; in response to an increment command, incrementing the address by the stride; and accessing the memory using the incremented address.

[0054] In Example 2, the subject matter of Example 1, wherein the accessing of the memory using the address comprises writing data to the address.

[0055] In Example 3, the subject matter of Example 2, wherein the operations further comprise: receiving the data using the same wires used to receive the address.

[0056] In Example 4, the subject matter of Examples 1-3, wherein the operations further comprise: receiving a second stride; in response to a second increment command, further incrementing the address by the second stride; and accessing the memory using the further incremented address.

[0057] In Example 5, the subject matter of Examples 1-4, wherein the receiving of the stride uses the same wires as the receiving of the address.

[0058] In Example 6, the subject matter of Examples 1-5, wherein the logic block is coupled to a switch box of a field programmable gate array (FPGA).

[0059] In Example 7, the subject matter of Examples 1-6, wherein the logic block comprises a plurality of look-up tables (LUTs) and an arithmetic logic unit (ALU).

[0060] Example 8 is a method comprising: receiving, by a memory block comprising a logic block and a memory, an address and a stride; accessing, by the logic block, the memory using the address; in response to an increment command, incrementing, by the logic block, the address by the stride; and accessing, by the logic block, the memory using the incremented address.

[0061] In Example 9, the subject matter of Example 8, wherein the accessing of the memory using the address comprises writing data to the address.

[0062] In Example 10, the subject matter of Example 9 includes receiving, by the logic block, the data using the same wires used to receive the address.

[0063] In Example 11, the subject matter of Examples 8-10 includes receiving, by the logic block, a second stride; in response to a second increment command, further incrementing the address by the second stride; and accessing, by the logic block, the memory using the further incremented address.

[0064] In Example 12, the subject matter of Examples 8-11, wherein the receiving of the stride uses the same wires as the receiving of the address.

[0065] In Example 13, the subject matter of Examples 8-12, wherein the logic block is coupled to a switch box of a field programmable gate array (FPGA).

[0066] In Example 14, the subject matter of Examples 8-13, wherein the logic block comprises a plurality of look-up tables (LUTs) and an arithmetic logic unit (ALU).

[0067] Example 15 is a non-transitory machine-readable medium that comprises instructions that, when executed by one or more processors, cause the one or more processors to configure a logic

block of a memory block of a field programmable gate array (FPGA) to perform operations comprising: receiving an address and a stride; accessing a memory of the memory block using the address; in response to an increment command, incrementing the address by the stride; and accessing the memory using the incremented address.

[0068] In Example 16, the subject matter of Example 15, wherein the accessing of the memory using the address comprises writing data to the address.

[0069] In Example 17, the subject matter of Example 16, wherein the operations further comprise: receiving the data using the same wires used to receive the address.

[0070] In Example 18, the subject matter of Examples 15-17, wherein the operations further comprise: receiving a second stride; in response to a second increment command, further incrementing the address by the second stride; and accessing the memory using the further incremented address.

[0071] In Example 19, the subject matter of Examples 15-18, wherein the receiving of the stride uses the same wires as the receiving of the address.

[0072] In Example 20, the subject matter of Examples 15-19, wherein the logic block is coupled to a switch box of the FPGA.

[0073] Example 21 is an apparatus comprising means to implement any of Examples 1-20.

[0074] FIG. 7 is a block diagram illustrating an example of a machine **700** upon or by which one or more example process embodiments described herein may be implemented or controlled. In alternative embodiments, the machine **700** may operate as a standalone device or may be connected (e.g., networked) to other machines. In a networked deployment, the machine **700** may operate in the capacity of a server machine, a client machine, or both in server-client network environments. In an example, the machine **700** may act as a peer machine in a peer-to-peer (P2P) (or other distributed) network environment. Further, while only a single machine **700** is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein, such as via cloud computing, software as a service (SaaS), or other computer cluster configurations.

[0075] Examples, as described herein, may include, or may operate by, logic, a number of components, or mechanisms. Circuitry is a collection of circuits implemented in tangible entities that include hardware (e.g., simple circuits, gates, logic). Circuitry membership may be flexible over time and underlying hardware variability. Circuitries include members that may, alone or in combination, perform specified operations when operating. In an example, hardware of the circuitry may be immutably designed to carry out a specific operation (e.g., hardwired). In an example, the hardware of the circuitry may include variably connected physical components (e.g., execution units, transistors, simple circuits) including a computer-readable medium physically modified (e.g., magnetically, electrically, by moveable placement of invariant massed particles) to encode instructions of the specific operation. In connecting the physical components, the underlying electrical properties of a hardware constituent are changed (for example, from an insulator to a conductor or vice versa). The instructions enable embedded hardware (e.g., the execution units or a loading mechanism) to create members of the circuitry in hardware via the variable connections to carry out portions of the specific operation when in operation. Accordingly, the computer-readable medium is communicatively coupled to the other components of the circuitry when the device is operating. In an example, any of the physical components may be used in more than one member of more than one circuitry. For example, under operation, execution units may be used in a first circuit of a first circuitry at one point in time and reused by a second circuit in the first circuitry, or by a third circuit in a second circuitry, at a different time.

[0076] The machine (e.g., computer system) **700** may include a hardware processor **702** (e.g., a central processing unit (CPU), a hardware processor core, or any combination thereof), a graphics processing unit (GPU) **703**, a main memory **704**, and a static memory **706**, some or all of which

may communicate with each other via an interlink (e.g., bus) **708**. The machine **700** may further include a display device **710**, an alphanumeric input device **712** (e.g., a keyboard), and a user interface (UI) navigation device **714** (e.g., a mouse). In an example, the display device **710**, alphanumeric input device **712**, and UI navigation device **714** may be a touch screen display. The machine **700** may additionally include a mass storage device (e.g., drive unit) **716**, a signal generation device **718** (e.g., a speaker), a network interface device **720**, and one or more sensors **721**, such as a Global Positioning System (GPS) sensor, compass, accelerometer, or another sensor. The machine **700** may include an output controller **728**, such as a serial (e.g., universal serial bus (USB)), parallel, or other wired or wireless (e.g., infrared (IR), near field communication (NFC)) connection to communicate with or control one or more peripheral devices (e.g., a printer, card reader).

[0077] The mass storage device **716** may include a machine-readable medium **722** on which is stored one or more sets of data structures or instructions **724** (e.g., software) embodying or utilized by any one or more of the techniques or functions described herein. The instructions **724** may also reside, completely or at least partially, within the main memory **704**, within the static memory **706**, within the hardware processor **702**, or within the GPU **703** during execution thereof by the machine **700**. In an example, one or any combination of the hardware processor **702**, the GPU **703**, the main memory **704**, the static memory **706**, or the mass storage device **716** may constitute machine-readable media.

[0078] While the machine-readable medium **722** is illustrated as a single medium, the term “machine-readable medium” may include a single medium, or multiple media, (e.g., a centralized or distributed database, and/or associated caches and servers) configured to store the one or more instructions **724**.

[0079] The term “machine-readable medium” may include any medium that is capable of storing, encoding, or carrying instructions **724** for execution by the machine **700** and that cause the machine **700** to perform any one or more of the techniques of the present disclosure, or that is capable of storing, encoding, or carrying data structures used by or associated with such instructions **724**. Non-limiting machine-readable medium examples may include solid-state memories, and optical and magnetic media. In an example, a massed machine-readable medium comprises a machine-readable medium **722** with a plurality of particles having invariant (e.g., rest) mass. Accordingly, massed machine-readable media are not transitory propagating signals. Specific examples of massed machine-readable media may include non-volatile memory, such as semiconductor memory devices (e.g., Electrically Programmable Read-Only Memory (EPROM), Electrically Erasable Programmable Read-Only Memory (EEPROM)) and flash memory devices; magnetic disks, such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks.

[0080] The instructions **724** may further be transmitted or received over a communications network **726** using a transmission medium via the network interface device **720**.

[0081] Throughout this specification, plural instances may implement components, operations, or structures described as a single instance. Although individual operations of one or more methods are illustrated and described as separate operations, one or more of the individual operations may be performed concurrently, and nothing requires that the operations be performed in the order illustrated. Structures and functionality presented as separate components in example configurations may be implemented as a combined structure or component. Similarly, structures and functionality presented as a single component may be implemented as separate components. These and other variations, modifications, additions, and improvements fall within the scope of the subject matter herein.

[0082] The embodiments illustrated herein are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed. Other embodiments may be used and derived therefrom, such that structural and logical substitutions and changes may be made without

departing from the scope of this disclosure. The Detailed Description, therefore, is not to be taken in a limiting sense, and the scope of various embodiments is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

[0083] As used herein, the term “or” may be construed in either an inclusive or exclusive sense. Moreover, plural instances may be provided for resources, operations, or structures described herein as a single instance. Additionally, boundaries between various resources, operations, modules, engines, and data stores are somewhat arbitrary, and particular operations are illustrated in a context of specific illustrative configurations. Other allocations of functionality are envisioned and may fall within a scope of various embodiments of the present disclosure. In general, structures and functionality presented as separate resources in the example configurations may be implemented as a combined structure or resource. Similarly, structures and functionality presented as a single resource may be implemented as separate resources. These and other variations, modifications, additions, and improvements fall within a scope of embodiments of the present disclosure as represented by the appended claims. The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A memory block comprising: a memory; and a logic block coupled to the memory and configured to perform operations comprising: sequentially receiving an address and a stride using the same wires; accessing the memory using the address; in response to an increment command, incrementing the address by the stride; and accessing the memory using the incremented address.
2. The memory block of claim 1, wherein the accessing of the memory using the address comprises writing data to the address.
3. The memory block of claim 2, wherein the operations further comprise: receiving the data using the same wires used to receive the address.
4. The memory block of claim 1, wherein the operations further comprise: receiving a second stride; in response to a second increment command, further incrementing the address by the second stride; and accessing the memory using the further incremented address.
5. The memory block of claim 1, wherein the receiving of the stride uses the same wires as the receiving of the address.
6. The memory block of claim 1, wherein the logic block is coupled to a switch box of a field programmable gate array (FPGA).
7. The memory block of claim 1, wherein the logic block comprises a plurality of look-up tables (LUTs) and an arithmetic logic unit (ALU).
8. A method comprising: receiving, by a memory block comprising a logic block and a memory, an address and a stride; accessing, by the logic block, the memory using the address; in response to an increment command, incrementing, by the logic block, the address by the stride; and accessing, by the logic block, the memory using the incremented address.
9. The method of claim 8, wherein the accessing of the memory using the address comprises writing data to the address.
10. The method of claim 9, further comprising: receiving, by the logic block, the data using the same wires used to receive the address.
11. The method of claim 8, further comprising: receiving, by the logic block, a second stride; in response to a second increment command, further incrementing the address by the second stride; and accessing, by the logic block, the memory using the further incremented address.
12. The method of claim 8, wherein the receiving of the stride uses the same wires as the receiving of the address.
13. The method of claim 8, wherein the logic block is coupled to a switch box of a field programmable gate array (FPGA).

- 14.** The method of claim 8, wherein the logic block comprises a plurality of look-up tables (LUTs) and an arithmetic logic unit (ALU).
- 15.** A non-transitory machine-readable medium that comprises instructions that, when executed by one or more processors, cause the one or more processors to configure a logic block of a memory block of a field programmable gate array (FPGA) to perform operations comprising: receiving an address and a stride; accessing a memory of the memory block using the address; in response to an increment command, incrementing the address by the stride; and accessing the memory using the incremented address.
- 16.** The non-transitory machine-readable medium of claim 15, wherein the accessing of the memory using the address comprises writing data to the address.
- 17.** The non-transitory machine-readable medium of claim 16, wherein the operations further comprise: receiving the data using the same wires used to receive the address.
- 18.** The non-transitory machine-readable medium of claim 15, wherein the operations further comprise: receiving a second stride; in response to a second increment command, further incrementing the address by the second stride; and accessing the memory using the further incremented address.
- 19.** The non-transitory machine-readable medium of claim 15, wherein the receiving of the stride uses the same wires as the receiving of the address.
- 20.** The non-transitory machine-readable medium of claim 15, wherein the logic block is coupled to a switch box of the FPGA.
-