

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265071

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Ananthesweran; Krishnan et al.

DEPLOYMENT ORCHESTRATOR IN A CLOUD BASED SERVICE FOR CUSTOMER ENTERPRISES

Abstract

A deployment orchestrator is provided that manages package deployments at different hierarchical levels. Each hierarchical level is associated with a particular type of resource object. The deployment orchestrator creates different of resource objects, each associated with a different hierarchical level and updates instances of the different resource objects based on information related to a package that is to be deployed. The deployment orchestrator performs processing associated with deploying the package at the hierarchical level based on information stored in the instances of the resource objects associated with the hierarchical level e.g., information related to a package that is to be deployed.

Inventors: Ananthesweran; Krishnan (San Jose, CA), Pinnapareddy; Nishanth Reddy (Rancho Cordova, CA), Zhan; Hao (Cupertino, CA), Chawla; Gaurav (San Jose, CA)

Applicant: Cisco Technology, Inc. (San Jose, CA)

Family ID: 1000008576630

Appl. No.: 19/198426

Filed: May 05, 2025

Related U.S. Application Data

parent US continuation 18180662 20230308 parent-grant-document US 12293179 child US 19198426

Publication Classification

Int. Cl.: G06F8/65 (20180101)

Background/Summary

RELATED APPLICATIONS [0001] This application is continuation of U.S. Non-Provisional application Ser. No. 18/180,662, filed on Mar. 8, 2023, and titled “DEPLOYMENT ORCHESTRATOR IN A CLOUD BASED SERVICE FOR CUSTOMER ENTERPRISES,” which is hereby incorporated by reference in its/their entirety for all purposes.

FIELD

[0002] The present disclosure relates to a cloud-based service that manages resources in a customer enterprise.

BACKGROUND

[0003] Modern information technology (IT) infrastructures e.g., data centers, customer enterprise networks, etc., often comprise thousands of computer systems that may operate collectively to service requests from even larger numbers of remote computing devices. In operation, such large-scale IT infrastructures generate significant volumes of performance data and diagnostic information that needs to be analyzed to diagnose performance and security problems. Software is utilized to manage the collection, storage, and analysis of such diagnostic information and performance data. The software may include components across the network, including at the data sources, at assigned storage locations, etc. For some network operators, the computer systems in such large-scale infrastructures are managed and/or controlled directly (e.g., behind a firewall of the enterprise) by a group of authorized personnel (e.g., administrators), who are employee-agents of the enterprise. Such large-scale infrastructures, where the network operator itself administers the computer systems that make up the infrastructures, are referred to as customer managed environments or self-managed environments.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] Features, implementations, and advantages of the present disclosure are better understood when the following Detailed Description is read with reference to the accompanying drawings.

[0005] FIG. 1 is a high-level diagram of an enterprise deployment environment that is managed by a cloud-based enterprise service according to certain implementations.

[0006] FIG. 2 depicts a customer enterprise that is managed by a cloud-based enterprise service according to certain implementation.

[0007] FIG. 3 depicts another customer enterprise implementation that is managed by the cloud-based enterprise service.

[0008] FIG. 4 depicts a relational diagram illustrating objects and their respective attributes that are exchanged between a supervisor and a cloud-based enterprise configuration controller.

[0009] FIG. 5 depicts a high-level diagram illustrating resources created in an onboarding process of a customer.

[0010] FIG. 6 is a flow diagram illustrating a customer on-boarding process.

[0011] FIG. 7 illustrates a simplified flowchart depicting steps performed by a node enterprise supervisor, according to certain implementations.

[0012] FIG. 8 illustrates a simplified flowchart depicting steps performed by the cloud-based enterprise configuration controller according to certain implementations.

[0013] FIG. **9** illustrates a simplified flowchart depicting steps performed by a cloud-based enterprise service in updating one or more configuration parameters of a package on one or more nodes of a customer enterprise, according to certain implementations.

[0014] FIG. **10** illustrates a simplified flowchart depicting steps performed by a server/engine supervisor, according to certain implementations.

[0015] FIG. **11** illustrates another simplified flowchart depicting steps performed by the cloud-based enterprise configuration controller according to certain implementations.

[0016] FIG. **12** illustrates a simplified flowchart depicting steps performed by a server service deployed in a cloud environment according to certain implementations.

[0017] FIG. **13** depicts an exemplary swim diagram illustrating steps performed for creating a cluster of servers according to certain embodiments.

[0018] FIG. **14** depicts a swim diagram illustrating steps performed in provisioning a cluster of servers according to certain embodiments.

[0019] FIG. **15** depicts an exemplary schematic illustrating components of a deployment orchestrator according to certain embodiments.

[0020] FIG. **16** depicts an exemplary system diagram illustrating operations of a deployment orchestrator according to certain embodiments.

[0021] FIG. **17** depicts a swim diagram illustrating steps performed in deploying a package according to certain embodiments.

[0022] FIG. **18** depicts an exemplary state transition diagram corresponding to one hierarchical level of a plurality of hierarchical levels according to certain embodiments.

[0023] FIG. **19** depicts an exemplary state transition diagram corresponding to a rollback operation according to certain embodiments.

[0024] FIG. **20** depicts an exemplary state transition diagram corresponding to another hierarchical level of a plurality of hierarchical levels according to certain embodiments.

[0025] FIG. **21** is a block diagram illustrating an example computing environment that includes a data intake and query system.

[0026] FIG. **22** is a block diagram illustrating an example of an indexing system included in the data intake and query system

[0027] FIG. **23** is a block diagram illustrating an example of a search system included in the data intake and query system.

[0028] FIG. **24** illustrates an example of a customer managed network.

DETAILED DESCRIPTION

[0029] In the following description, for the purposes of explanation, specific details are set forth in order to provide a thorough understanding of certain implementations. However, it will be apparent that various implementations may be practiced without these specific details. The figures and description are not intended to be restrictive. Any implementation or design described herein as example is not necessarily to be construed as preferred or advantageous over other implementations or designs.

[0030] Entities of various types, such as companies, educational institutions, medical facilities, governmental departments, and private individuals, among other examples, operate large-scale information technology ecosystems for various purposes. An information technology ecosystem includes infrastructures of various kinds of computing resources including computer systems, servers, storage systems, network communication devices, or any other electronic resources. In such information technology ecosystems (also referred to herein as customer enterprises), the computing resources may be managed by authorized personnel of the enterprise (e.g., administrators). For instance, software may be deployed on the computing resources of the enterprise for the purpose of data collection, data storage, and data analytics. The authorized personnel of the enterprise may handle the management and administration of such a large-scale software deployment behind a firewall of the enterprise.

[0031] Computing environments, which can also be referred to as information technology environments, can include inter-networked, physical hardware devices, the software executing on the hardware devices, and the users of the hardware and software. As an example, an entity such as a school can operate a Local Area Network (LAN) that includes desktop computers, laptop computers, smart phones, and tablets connected to a physical and wireless network, where users correspond to teachers and students. In this example, the physical devices may be in buildings or a campus that is controlled by the school.

[0032] As another example, an entity such as a business can operate a Wide Area Network (WAN) that includes physical devices in multiple geographic locations where the offices of the business are located. In this example, the different offices can be inter-networked using a combination of public networks such as the Internet and private networks. As another example, an entity can operate a data center: a centralized location where computing resources are kept and maintained, and whose resources are accessible over a network. In this example, users associated with the entity that operates the data center can access the computing resources in the data center over public and/or private networks that may not be operated and controlled by the same entity. Alternatively, or additionally, the operator of the data center may provide the computing resources to users associated with other entities, for example on a subscription basis. In both of these examples, users may expect resources to be available on demand and without direct active management by the user, i.e., a resource delivery model often referred to as cloud computing.

[0033] Entities that operate computing environments need information about their computing resources. For example, an entity may need to know the operating status of the various computing resources in the entity's computing environment, so that the entity can administer the environment, including performing configuration and maintenance, performing repairs or replacements, provisioning additional resources, removing unused resources, or addressing issues that may arise during operation of the computing environment, among other examples. As another example, an entity can use information about a computing environment to identify and remediate security issues that may endanger the data, users, and/or equipment in the computing environment. As another example, an entity may be operating a computing environment for some purpose (e.g., to run an online store, to operate a bank, to manage a municipal railway, etc.) and information about the computing environment can aid the entity in understanding whether the computing environment is serving its purpose well.

[0034] To obtain information about its computing environments, an entity can use software designed for this purpose. This software can be deployed across a computing environment, such as, for example, on some or all of the various computing devices that make up the computing environment. In some cases, the entity may manage such a software deployment itself, for example by having employees assigned to manage the software deployment, who have the proper network authorization to do so. The management functions are handled within the entity's computing environment and within the security and access boundaries (e.g., the firewall) of the computing environment. Such an environment may be referred to as a self-managed environment.

[0035] In self-managed environments, a frequently encountered problem is that of scalability. As the IT ecosystem grows in size, managing deployment and administration of software in the individual nodes (i.e., compute resources of the environment) becomes increasingly difficult. As such, a new service with a new architecture is required that is capable of delivering software and configuration updates to customer environments in a seamless and scalable manner. As described herein, there is provided a cloud-based mechanism for delivering and managing software services to various customer environments. Specifically, a cloud-based service manages resources in one or more host machines included in a customer enterprise deployment. A supervisor that is instantiated in each node, communicates with a configuration controller included in the cloud-based service. The configuration controller provides a seamless manner to control and manage resources deployed on one or more host machines included in the customer enterprise deployment(s).

[0036] Turning now to FIG. 1, there is depicted a high-level diagram of an enterprise deployment environment that includes software that is managed by a cloud-based enterprise service according to certain implementations. As shown in FIG. 1, an enterprise deployment environment **101** includes a plurality of customer enterprises. For instance, the enterprise deployment environment **101** includes a first customer enterprise (i.e., customer A enterprise **103**), a second customer enterprise (i.e., customer B enterprise **113**), and a third customer enterprise (i.e., customer C enterprise **123**). Each of the first, second, and third customer enterprises include software that is managed by a cloud-based enterprise service **140** i.e., the first, second, and the third customer enterprises are customers of the cloud-based enterprise service. The cloud-based enterprise service **140** is also referred to herein as a cloud service **140** or a cloud service infrastructure **140**. The customer enterprises communicate with the cloud service **140** over a public network e.g., the Internet. It is appreciated that the enterprise deployment environment **101** may make use of the services offered by the cloud service **140** via a firewall **130**. The enterprise deployment environment **101** is shown to include three customer enterprises for the sake of illustration. However, this is in no way limiting the scope of the present disclosure, and it is appreciated that the enterprise deployment environment **101** may include any number of customer enterprise(s).

[0037] Each customer enterprise may include one or more nodes i.e., computing resources that execute one or more software packages. A software package (also referred to herein as simply a package) is a standalone software unit that is defined by an application binary and configuration setting(s) of the application. For instance, as shown in FIG. 1, customer A enterprise **103** includes one or more nodes that execute packages **106**, customer B enterprise **113** includes one or more nodes that execute packages **116**, and customer C enterprise **123** includes one or more nodes that execute packages **126**, respectively.

[0038] According to some implementations, each node within a customer enterprise includes an instance of a node enterprise supervisor. For example, as shown in FIG. 1, an enterprise node **104** (included in customer A enterprise **103**) comprises a node enterprise supervisor **105**, whereas enterprise node **114** (included in customer B enterprise **113**) comprises a node enterprise supervisor **115**, and enterprise node **124** (included in customer C enterprise **123**) comprises a node enterprise supervisor **125**. Each of the node enterprise supervisors corresponds to a client-side software (e.g., host process) that brokers communication between the cloud service and the individual package(s) that are executing on the node. Moreover, each supervisor may be programmed to manage: (a) downloading of a package from the cloud service **140**, and (b) executing and updating packages deployed on the node. In some implementations, a node enterprise supervisor can be uniquely identified by a self-generated identifier (i.e., ID) and can execute one or more packages on the node, where each package is associated with a unique configuration setting.

[0039] The cloud service **140** includes a package content delivery network **141**, a package release unit **143**, a cloud based enterprise configuration controller **145** (referred to herein as a configuration controller), a package service **147**, and a deployment orchestrator **149**. The package release unit **143** is programmed to publish a new version of a package and make the new version available for deployment. In some implementations, the new version of the package is released to the package content delivery network **141**. The node supervisors (associated with enterprise node(s)) communicate with the package content delivery network **141** to obtain the latest version of a particular package. The configuration controller **145** manages, monitors package instances, and orchestrates deployments of packages (and configuration settings of the packages) in the enterprise node(s). In other words, the configuration controller **145** is configured to ensure that the latest version of a package is up and running on an enterprise node. Additionally, it is noted that the cloud service cloud service **140** may be provided/executed in different geographical locations (i.e., regions) to service the enterprise deployments in the corresponding geographical locations.

[0040] The package service **147** provides domain-specific application interface(s) (i.e., APIs) in the cloud that may be used by packages for package configuration administration. The deployment

orchestrator **149** can be configured to push (via the configuration controller **145**) updates for a particular package(s) on all nodes of a particular enterprise i.e., on a global level. Moreover, the deployment orchestrator **149** can be utilized in conjunction with the package service **147** in incrementally updating packages on a set of one or more nodes. Details regarding the interactions of the cloud service **140** with the node enterprise supervisors is described next with reference to FIGS. **2** and **3**, respectively. For sake of simplicity, the interactions are described with regard to a single customer enterprise but are equally applicable to all customer enterprises included in the enterprise deployment environment **101** of FIG. **1**.

[0041] FIG. **2** depicts a customer enterprise that includes software that is managed by the cloud service according to certain implementations. Specifically, a customer enterprise **201** includes one or more enterprise nodes **203**, where packages executed on the respective one or more enterprise nodes **203** are managed by cloud service **240**. In some implementations, each node within the customer enterprise **201** comprises a node enterprise supervisor **204** that downloads (from the cloud service **240**) and executes the downloaded packages **205** on the node. It is noted that the cloud service **240** is secured from the customer enterprises via a firewall **230**. The node enterprise supervisor **204** communicates with the cloud service **240** over a public network such as the Internet. The cloud service **240** includes a package content delivery network **241**, a package release unit **243**, a configuration controller **245**, a package service **247**, and a deployment orchestrator **249**. It is appreciated that the functionality of the individual components of the cloud service **240** are similar to the functionality of the components of the cloud service **140** described above with reference to FIG. **1**.

[0042] In one implementation, the node enterprise supervisor **204** of FIG. **2** transmits a recurring heartbeat message to the configuration controller **245** included in the cloud service **240**. The heartbeat message sent from the node enterprise supervisor **204** to the configuration controller **245** announces the supervisor instance's existence, and includes metadata (e.g., first metadata) that identifies labels, a current configuration (i.e., a version of the package and/or configuration settings of the package) of a package that is executed on the node. Additionally, the first metadata may include information indicative of a metric associated with the node enterprise supervisor. For instance, the metric may correspond to a health metric of the node enterprise supervisor **204**, which contains time-series data-points that are used to capture health and other operational data of the node enterprise supervisor. Moreover, the frequency of transmitting the heartbeat messages may be in the order of a few minutes e.g., every ten minutes. In some implementations, the node enterprise supervisor **204** may utilize other communication mechanisms to enable the cloud service **240** to asynchronously initiate communication with the node enterprise supervisor **204** and thus reduce communication latency. For example, a long polling request-response channel or a bi-directional communication channel may be established between the node enterprise supervisor **204** and the configuration controller **245** (included in the cloud service **240**) to exchange information in a seamless manner.

[0043] The configuration controller **245** upon receiving a heartbeat message determines whether the package executed by the node enterprise supervisor is up to date. In other words, the configuration controller **245** determines whether the package executed by the node enterprise supervisor is to be updated. With regard to updating a package, the configuration controller **245** may determine whether a version of the package is to be updated and/or whether one or more configuration settings associated with the package are to be updated. It is appreciated that the configuration controller **245** may perform such a determination by comparing the metadata of the package that is included in the heartbeat message transmitted from the node enterprise supervisor to information (i.e., version and/or configuration settings) of a newer package, for example that is published to the package content delivery network **241**.

[0044] Responsive to receiving the heartbeat message, the configuration controller **245** transmits a configuration response message back to the node enterprise supervisor **204**. Upon successfully

determining that a current package (i.e., the package currently being executed on the node) is to be updated (e.g., version of the package and/or configuration setting(s) of the package have changed), the configuration controller **245** transmits a configuration response message back to the node enterprise supervisor **204**. The configuration response message includes metadata (e.g., second metadata) indicative of a newer version and/or configuration settings of the package, and a location on the package content delivery network **241** (e.g., an URL) from where the node enterprise supervisor **204** may download the required updated package. Upon receiving the above-described configuration response message, the node enterprise supervisor **204** communicates with the package content delivery network **241** to download the updated package and executes the downloaded package on the node. In some implementations, the configuration response message includes configuration settings that may be classified as one of a secret configuration setting (e.g., configuration settings used for authentication purposes) or a non-secret configuration setting. As secret configuration settings are confidential in nature, they may be transmitted from the configuration controller **245** to the node enterprise supervisor **204** in a manner that is different than the transmission of non-secret configuration settings. For instance, the non-secret configuration settings may be included in the metadata of the configuration response message. In contrast, the secret configuration settings may be included in a separate data object included in the heartbeat response messages. In this manner, the secret configuration setting may be processed separately from the non-secret configuration settings, while ensuring that the secret configuration settings is not persisted to a storage location that is not properly protected from unauthorized access. In contrast, when the configuration controller **245** determines that the current package is not required to be updated i.e., the version of the package and the one or more configuration settings of the package are up to date, then the configuration controller **245** transmits the configuration response message to the node enterprise supervisor simply indicating that no updates are required.

[0045] In another implementation, an administrator of the customer enterprise **201** may edit one or more configuration settings of a package that are to be executed across several nodes within the customer enterprise **201**. For example, the administrator may utilize an administrator user interface (UI) **207** to transmit, one or more edited configuration settings to the package service **247**. The package service **247** upon receiving the edited one or more configuration settings, publishes the edited one or more configuration settings to the configuration controller **245**. In turn, the configuration controller **245** transmits the edited one or more configuration settings to the respective node enterprise supervisors (of the nodes included in the customer enterprise **201**) in heartbeat response messages that are transmitted to the respective nodes upon the configuration controller **245** receiving heartbeat messages from the nodes.

[0046] In some implementations, the cloud service **240** may utilize the deployment orchestrator **249** in conjunction with the package service **247**. For instance, in the case of updating the one or more configuration settings on one or more nodes of the customer enterprise **201**, the deployment orchestrator **249** may be programmed to update the configuration settings on the set of one or more nodes in an incremental fashion. Specifically, the deployment orchestrator **249** may instruct the configuration controller **245** to update the configuration setting on a first node and verify whether a heartbeat message transmitted from the first node is okay, before proceeding to update the configuration setting on subsequent node. Moreover, the deployment orchestrator **249** may also be configured to perform a rollback operation i.e., the deployment orchestrator **249** may instruct the configuration controller **245** to revert back to a previous version of the package in case of any issues being faced while updating the set of one or more nodes.

[0047] The cloud service can be utilized to manage resources in customer enterprises in a variety of applications. For instance, customer enterprises can include a server/engine referred to as an edge routing device or edge server. In such an example, the server can act as an intermediate forwarder between different sources. Specifically, in an enterprise setting, customer(s) perform data processing operations such as data filtering, transform operations, data reductions, etc. The

processed data is desired to be stored in different data storages. For example, a first portion of the data may be stored in a first data storage at a first geographical location, and a second portion of the data may be stored in a second data storage at a second geographical location (different than the first geographical location). In such a scenario, the server may be used as an intermediate forwarder to route specific data to a desired storage devices.

[0048] The server can be characterized by configuration settings that enable the server to perform the desired tasks. In such an enterprise setting, there are a plurality of such servers that perform the above-described data processing operations. Typically, a user with sufficient privileges (e.g., an administrator) would have to manually set the configuration settings on each server. When a customer enterprise includes tens or even only hundreds of such servers, manually configuring each one is not feasible. In what follows, there is described with reference to FIG. 3, a mechanism to utilize the cloud service to manage such servers in a customer's enterprise.

[0049] FIG. 3 depicts another customer enterprise implementation that includes software that is managed by the cloud service. An example customer enterprise **301** includes one or more servers **303** (i.e., edge servers that act as intermediate forwarders between data sources). For instance, an edge server **305** in the customer enterprise forwards data to different data sources. In some implementations, each server within the customer enterprise **301** comprises a server supervisor **304** that downloads (from a cloud service **340**) and installs packages on the server. It is noted that the cloud service **340** is secured from the customer enterprise **301** via a firewall **330**. The server supervisor **304** communicates with the cloud service **340** over a public network such as the Internet.

[0050] The cloud service **340** includes a package content delivery network **341**, a server release unit **343**, a configuration controller **345**, a server configuration service **347**, and a deployment orchestrator **349**. It is appreciated that the functionality of the individual components of the cloud service **340** are similar to the functionality of the components of the cloud service **240** described above with reference to FIG. 2.

[0051] In one implementation, the server supervisor **304** of FIG. 3 transmits a recurring heartbeat message to the configuration controller **345** included in the cloud service **340**. The heartbeat message sent from the server supervisor **304** to the configuration controller **345** announces the supervisor instance's existence, identifies label metadata, and current configuration of a package that is executed on the server. Specifically, each heartbeat message may include metadata (e.g., first metadata) that includes the above stated information as well as includes information indicative of a metric associated with the server supervisor. For instance, the metric may correspond to a health metric of the server supervisor **304**, which contains time-series data-points that are used to capture health and other operational data of the server supervisor. Moreover, the frequency of transmitting the heartbeat messages may be in the order of a few minutes e.g., every ten minutes.

[0052] The configuration controller **345** upon receiving a heartbeat message determines whether the package executed by the server supervisor is up to date. In other words, the configuration controller **345** determines whether the package executed by the server supervisor is to be updated. With regard to updating a package, the configuration controller **345** may determine whether a version of the package is to be updated and/or whether one or more configuration settings associated with the package are to be updated. It is appreciated that the configuration controller **345** may perform such a determination by comparing the metadata of the package that is included in the heartbeat message transmitted from the server supervisor to information (i.e., version and/or configuration settings) of a newer package that is, for example, published to the package content delivery network **341** via the server release unit **343**.

[0053] Responsive to receiving the heartbeat message, the configuration controller **345** transmits a configuration response message back to the server supervisor **304**. Upon successfully determining that a current package (i.e., the package currently being executed on the server) is to be updated (e.g., version of the package and/or configuration setting(s) of the package have changed), the

configuration controller **345** transmits a configuration response message back to the server supervisor **304**. The configuration response message includes metadata (e.g., second metadata) indicative of a newer version and/or configuration settings of the package, and a location on the package content delivery network **341** (e.g., an URL) from where the server supervisor **304** may download the required updated package. Upon receiving the above-described configuration response message, the server supervisor **304** communicates with the package content delivery network **341** to download the updated package and executes the downloaded package on the server **305**. In contrast, when the configuration controller **345** determines that the current package is not required to be updated i.e., the version of the package and the one or more configuration settings of the package are up to date, then the configuration controller **345** transmits the configuration response message to the server supervisor simply indicating that no updates are required.

[0054] In another implementation, an administrator of the customer enterprise **301** may edit one or more configuration settings of a package (executed on a server) across several server deployments within the customer enterprise **301**. For example, the administrator may utilize an administrator user interface (UI) **307** to transmit one or more edited configuration settings to the server configuration service **347**. The server configuration service **247**, upon receiving the edited one or more configuration settings, publishes the edited one or more configuration settings to the configuration controller **345**. In turn, the configuration controller **345** transmits the edited one or more configuration settings to the respective server supervisors (of the servers included in the customer enterprise **301**) in heartbeat response messages that are transmitted to the respective server supervisors in response to the configuration controller **345** receiving heartbeat messages from the server supervisors. Upon receiving the updated configuration setting(s), the server supervisors execute the configuration settings on the respective servers. In this manner, the cloud service **340** provides a seamless mechanism to manage edge servers in a customer enterprise.

[0055] Turning to FIG. 4, there is depicted an example of a relational diagram illustrating objects and their respective attributes that are exchanged between a supervisor and a cloud based enterprise configuration controller. The heartbeat message **401** is a recurring message sent from a supervisor to the configuration controller to announce the supervisor instance's existence. As shown in FIG. 4, the heartbeat message **401** includes a supervisor identifier (ID), labels that include metadata to facilitate grouping, a configuration hash that corresponds to a current configuration state, and metrics such as health metrics of the node.

[0056] The package configuration **405** corresponds to an object that is transmitted by the configuration controller in response to receiving the heartbeat message. The heartbeat response message indicates if the supervisor instance needs to update its software packages or their settings. In some implementations the package configuration object includes information pertaining to a name of the configuration, a version of the configuration, a URL where the configuration may be downloaded from (e.g., from content delivery network network), configuration settings and a timestamp indicating as to when the settings were updated. A supervisor represents a 'host' or 'lifecycle manager' that runs in the client environment and hosts one or more downloadable packages and their configuration. Each supervisor instance heartbeats continuously to the configuration controller to announce its initial presence and then its current configuration and health on an ongoing basis.

[0057] As shown in FIG. 4, the supervisor object **403** includes an ID of the supervisor and a configuration hash that corresponds to a configuration identifier unique to the current configuration of the supervisor instance. In some implementations, the configuration hash specifically is the hash of a package configurations field. It is appreciated that this field can be used to efficiently compare if there are any deltas (i.e., changes) in a current and target configuration. Furthermore, the supervisor object **403** includes a package configuration list that corresponds to a list of package versions and configurations that the supervisor is configured to use. Further, the metric object contains a time-series data points that are used to capture health and other operational data about a

supervisor (i.e., node enterprise supervisor). The supervisor sends metric data in its heartbeat message. The metric object **407** includes a name of the metric, a time stamp when the metric was observed, and the observed value of the metric.

[0058] FIG. 5 depicts a high-level diagram illustrating resources generated in an onboarding process of a customer. The process of onboarding a customer broadly relates to the generation of a tenancy (e.g., an account) for the customer so that the customer can avail services offered by a cloud-based enterprise service. To scale and to provide a modern cloud experience, enterprise customers should be able to self-service onboard. One difficulty with self-service onboarding is establishing trust with actors that have no cloud identity, mitigating any threats to unauthenticated onboarding APIs while still providing a low-friction user experience. More generally, onboarding entails establishing the right cloud resources (tenant, user accounts, client secrets, and “trust”) so that enterprise users and client processes can access authenticated cloud APIs.

[0059] In some implementations, the process of onboarding a customer involves the following steps: (1) generating a new a cloud tenancy, (2) installing a supervisor e.g., a node enterprise supervisor, (3) establishing authenticated service-to-service connections for node enterprise supervisors to the cloud based enterprise service, (4) validating customer licenses and subscriptions, and (5) establishing user identities for user administrators.

[0060] FIG. 5 depicts resources that are generated in a customer enterprise **501** and in a cloud-based enterprise service (cloud service) **551** as a result of the onboarding process. A node enterprise supervisor **504** is instantiated with respect to each node included in the customer enterprise. A user e.g., an administrator of the customer enterprise **501** may access the cloud-based enterprise service **551** via a browser **503**. The customer enterprise is associated with one or more licenses **506**. A service principal (i.e., an identity) is provisioned for an application executed in the customer enterprise to perform a service-to-service authentication to cloud services. In other words, the service principal functions as an identity of the application and may define parameters including what resource(s) can be accessed by the application. As shown in FIG. 5, the service principal **505** on the customer enterprise side is associated with a principal ID and a private key, whereas the corresponding service principal **557** included in the cloud service **551** includes the principal ID and a public key. It is noted that the public-private key pair is generated on the client side (i.e., the customer enterprise **501**) and only the public key is transmitted to the cloud service **551**.

[0061] A configuration object **555** is generated in the cloud service for each customer enterprise deployment. The configuration object **555** is used for bootstrapping customer onboarding, and includes parameters such as a name, a region, and a hash of a license that associates the configuration object **555** to a particular customer enterprise. An identity (e.g., user principal **559**) is created in the cloud service **551** for an administrator of the customer enterprise. The user principal is associated with an information related to the administrator e.g., an email address and password combination that is set by the administrator in the region. A tenancy **553** is created in the cloud service **551** that includes a user-friendly name of the tenancy and a region in which the tenancy included. It is appreciated that the tenancy **553** is related to the customer enterprise deployment **501** via a license key. Further, a tenant administrator group **563** is created to include the identities of one or more administrative accounts (associated with one or more administrators of the customer enterprise). A tenant membership **565** includes a list of all users and administrators of the customer enterprise that have access to the cloud service **551**. Furthermore, it is noted that the service principal **557** included in the cloud service **551** is a tenant member (i.e., associated with node enterprise supervisor access group **561**) and has access (i.e., can communicate with) the node enterprise supervisor **504**. It is appreciated that the above-described onboarding process is in no way limiting the scope of the present invention. Other mechanisms may be implemented to onboard a customer. For instance, a customer may obtain an ‘onboarding code’ to be successfully on boarded to the cloud service **551**.

[0062] FIG. 6 is a flow diagram illustrating a customer on-boarding process. For instance, FIG. 6

illustrates messages exchanged between a node in a customer enterprise **601** and a cloud-based service (cloud service) **610**. It is appreciated that the node in the customer enterprise may execute a package i.e., an application **603**, and be associated with an enterprise administrator **605**.

[0063] In step **1**, a configuration object is pre-provisioned in the cloud service **610**. In some implementations, there is a single configuration object provisioned per customer enterprise. Further, the configuration object (e.g., configuration object **555** as shown in FIG. 5) is characterized by a plurality of parameters such as a name, a region, and a license hash (associated with a customer enterprise license).

[0064] In step **2**, the application **603** executed on the node in the customer enterprise obtains a copy of the license (e.g., from the administrator). In step **3**, the application **603** makes an unauthenticated call (i.e., a query) to the cloud service **610** to check whether a tenancy is created (by the cloud service) in the cloud environment for the customer enterprise **601**. As the customer enterprise **601** has not yet been on-boarded to utilize the services of the cloud service **610**, in step **4**, the cloud service **610** transmits a response message back to the application **603** indicating that a customer tenancy has not yet been created for the customer enterprise.

[0065] In step **5**, the application **603** generates a key pair including a public key and a private key. The private key is stored locally in the node e.g., in a password configuration file. In step **6**, a request is transmitted from the node in the customer enterprise **601** to the cloud service **610** to create a tenancy for the customer. In some implementations, the request includes the license (obtained in step **2**) and the public key (generated in step **5**). In step **7**, the cloud service **610** validates the license based on the pre-provisioned configuration object (generated in step **1**). Upon successful validation, the cloud service **610** creates a tenancy for the customer enterprise **601** and adds the public key to a service principal (e.g., service principal **557** of FIG. 5).

[0066] In step **8**, the application **603** of FIG. 6 may poll the cloud service **610** to obtain status with respect to the request made in step **6**. In step **9**, the cloud service **610** responds with a message indicating that the tenancy is created for the customer enterprise, and also provides a principal identifier (i.e., principal ID) for the customer enterprise. The principal ID corresponds to an identifier that is used by the customer enterprise in future communications with the cloud service **610**. In step **10**, the application **603** stores the principal ID obtained in step **9**. In step **11**, the application transmits a request to authenticate the service principal on the customer enterprise side (e.g., service principal **505** of FIG. 5). In step **12**, the cloud service **610** transmits a token (referred to herein as SP token) which corresponds to a token associated with the service principal. The node in the customer enterprise **601** utilizes the obtained SP token in heartbeat messages that are transmitted periodically by the node to report, for instance, a version, configuration settings, metric health of the node, etc., to the cloud service **610**. In this manner, the customer enterprise **601** is on boarded to the cloud service **610**.

[0067] FIG. 7 illustrates a simplified flowchart **700** depicting steps performed by a node enterprise supervisor, according to certain implementations. The processing depicted in FIG. 7 may be implemented in software (e.g., code, instructions, program) executed by one or more processing units (e.g., processors, cores) of the respective systems, hardware, or combinations thereof. The software may be stored on a non-transitory storage medium (e.g., on a memory device). The method presented in FIG. 7 and described below is intended to be illustrative and non-limiting. Although FIG. 7 depicts the various processing steps occurring in a particular sequence or order, this is not intended to be limiting. In certain alternative implementations, the steps may be performed in some different order or some steps may also be performed in parallel.

[0068] The process commences in step **702**, where a node enterprise supervisor, e.g., the node enterprise supervisor, **204** of FIG. 2, transmits a heartbeat message to a configuration controller (e.g., configuration controller **245** of FIG. 2). The heartbeat message includes first metadata indicative of a current configuration of a package that is being executed on a node. It is noted that the current configuration may include include information pertaining to a version of the package

and/or one or more configuration settings of the package that is executed on the node. The heartbeat message may correspond to heartbeat message **401** of FIG. **4**.

[0069] In step **704**, the node enterprise supervisor, receives a heartbeat response message from the configuration controller. For instance, the heartbeat response message may correspond to the package configuration message **405** of FIG. **4**. The heartbeat response message includes second metadata indicative of whether the current configuration of the package is to be updated.

Thereafter, the process moves to step **706**, where responsive to the current configuration of the package requiring to be updated, the node enterprise supervisor, transmits a request message to a content delivery network (e.g., package content delivery network **241** of FIG. **2**) of the cloud service to download the updated version of the package. It is appreciated that in case of requiring an update of the version of the package, the node enterprise supervisor, downloads the updated version from the content delivery network, whereas in case of requiring an update to the one or more configuration settings of the package, the node enterprise supervisor may obtain (in some implementations) the updated configuration settings in the heartbeat response message received from the configuration controller. The process then moves to step **708**, where the node enterprise supervisor executes the downloaded version of the package and/or one or more updated configuration settings of the package in the node.

[0070] FIG. **8** illustrates a simplified flowchart **800** depicting steps performed by the cloud-based enterprise configuration controller (e.g., the configuration controller **245** of FIG. **2**) according to certain implementations. The processing depicted in FIG. **8** may be implemented in software (e.g., code, instructions, program) executed by one or more processing units (e.g., processors, cores) of the respective systems, hardware, or combinations thereof. The software may be stored on a non-transitory storage medium (e.g., on a memory device). The method presented in FIG. **8** and described below is intended to be illustrative and non-limiting. Although FIG. **8** depicts the various processing steps occurring in a particular sequence or order, this is not intended to be limiting. In certain alternative implementations, the steps may be performed in some different order or some steps may also be performed in parallel.

[0071] The process commences in step **802**, where the configuration controller receives a heartbeat message from a node enterprise supervisor. The heartbeat message includes a first metadata indicative of a current configuration of a package that is being executed on a node. Note that the current configuration may include information pertaining to a version of the package and/or one or more configuration settings of the package that is executed on the node by the node enterprise supervisor.

[0072] The process then moves to step **804**, where the configuration controller determines whether the current configuration of the package is to be updated. In some implementations, the configuration controller performs such a determination by comparing the first metadata included in the heartbeat message (transmitted from the node enterprise supervisor) to information (i.e., version and/or configuration settings) of a newer package that is published to a package content delivery network. Further, in step **806**, in response to successfully determining that current configuration of the package (i.e., the version of the package and/or one or more configuration settings of the package) is to be updated, the configuration controller transmits a heartbeat response to the node enterprise supervisor. The heartbeat response message includes second metadata indicative of a new configuration of the package that is to be downloaded by the node enterprise supervisor. It is appreciated that the node enterprise supervisor upon receiving the heartbeat response message downloads the updated package from the content delivery network. Moreover, it is appreciated that in case the configuration controller determines that the current package is not required to be updated i.e., the version of the package and the one or more configuration settings of the package are up to date, then the configuration controller transmits the configuration response message to the node enterprise supervisor simply indicating that no updates are required.

[0073] FIG. **9** illustrates a simplified flowchart **900** depicting steps performed by a cloud-based

enterprise service in updating one or more configuration parameters of a package on one or more nodes of a customer entries, according to certain implementations. In some implementations, the processing depicted in FIG. 9 may be implemented by the package service module 247 and the configuration controller 245 (i.e., components of the cloud-based enterprise service) depicted in FIG. 2. The processing depicted in FIG. 9 may be implemented in software (e.g., code, instructions, program) executed by one or more processing units (e.g., processors, cores) of the respective systems, hardware, or combinations thereof. The software may be stored on a non-transitory storage medium (e.g., on a memory device). The method presented in FIG. 9 and described below is intended to be illustrative and non-limiting. Although FIG. 9 depicts the various processing steps occurring in a particular sequence or order, this is not intended to be limiting. In certain alternative implementations, the steps may be performed in some different order or some steps may also be performed in parallel.

[0074] The process commences in step 902, where the package service (deployed in a cloud environment) receives a request from a user interface (e.g., admin UI 207 of FIG. 2) to apply updated, one or more configuration settings to a package that is executed on one or more nodes in a customer enterprise. The process of FIG. 9 then moves to step 904, where the package service publishes the updated one or more configuration settings to the configuration controller. Thereafter, the process moves to step 906, where the configuration controller transmits, in response to receiving a heartbeat message from each of the one or more nodes, a heartbeat response message including the updated configuration settings. In some implementations, the cloud service may utilize a deployment orchestrator (e.g., deployment orchestrator 249 of FIG. 2) to update the configuration settings on the one or more nodes in an incremental fashion. Specifically, the deployment orchestrator may instruct the configuration controller to update the configuration setting on a first node and verify whether a heartbeat message transmitted from the first node is okay, before proceeding to update the configuration setting on subsequent node.

[0075] FIG. 10 illustrates a simplified flowchart 1000 depicting steps performed by a server/engine supervisor, according to certain implementations. The processing depicted in FIG. 10 may be implemented in software (e.g., code, instructions, program) executed by one or more processing units (e.g., processors, cores) of the respective systems, hardware, or combinations thereof. The software may be stored on a non-transitory storage medium (e.g., on a memory device). The method presented in FIG. 10 and described below is intended to be illustrative and non-limiting. Although FIG. 10 depicts the various processing steps occurring in a particular sequence or order, this is not intended to be limiting. In certain alternative implementations, the steps may be performed in some different order or some steps may also be performed in parallel.

[0076] The process commences in step 1002, where a server supervisor e.g., server supervisor 304 of FIG. 3, transmits a heartbeat message to a configuration controller (e.g., configuration controller 345 of FIG. 3). The heartbeat message includes first metadata indicative of a current configuration of a server. It is noted that the current configuration may include information pertaining to a version of a package (executed on the server) and/or one or more configuration settings of the package.

[0077] In step 1004 of FIG. 10, the server supervisor receives a heartbeat response message from the configuration controller. For instance, the heartbeat response message may correspond to the package configuration message 405 of FIG. 4. The heartbeat response message includes second metadata indicative of whether the current configuration of the server is to be updated. Thereafter, the process of FIG. 10 moves to step 1006, where responsive to the current configuration of the server requiring to be updated, the server supervisor transmits a request message to a content delivery network (e.g., package content delivery network 341 of FIG. 3) of the cloud service 340 to download an updated configuration of the server (e.g., an updated version of the package executed on the server and/or updated configuration settings of the server). It is appreciated that in case of requiring an update of the version of the package, the server supervisor downloads the updated

version from the content delivery network, whereas in case of requiring an update to the one or more configuration settings of the package, the server supervisor may obtain (in some implementations) the updated configuration settings in the heartbeat response message received from the configuration controller. The process then moves to step **1008**, where the server supervisor executes the downloaded version of the package and/or one or more updated configuration settings on the server.

[0078] FIG. **11** illustrates another simplified flowchart **1100** depicting steps performed by the cloud-based enterprise configuration controller (e.g., the configuration controller **345** of FIG. **3**) according to certain implementations. The processing depicted in FIG. **11** may be implemented in software (e.g., code, instructions, program) executed by one or more processing units (e.g., processors, cores) of the respective systems, hardware, or combinations thereof. The software may be stored on a non-transitory storage medium (e.g., on a memory device). The method presented in FIG. **11** and described below is intended to be illustrative and non-limiting. Although FIG. **11** depicts the various processing steps occurring in a particular sequence or order, this is not intended to be limiting. In certain alternative implementations, the steps may be performed in some different order or some steps may also be performed in parallel.

[0079] The process commences in step **1102**, where the configuration controller receives a heartbeat message from a server supervisor. The heartbeat message includes a first metadata indicative of a current configuration of a server. Note that the current configuration may include information pertaining to a version of a package (executed on the server) and/or one or more configuration settings of the package.

[0080] The process then moves to step **1104**, where the configuration controller determines whether the current configuration of the server is to be updated. In some implementations, the configuration controller performs such a determination by comparing the first metadata included in the heartbeat message (transmitted from the server supervisor) to information (i.e., version and/or configuration settings) of a newer package that is published to a package content delivery network. Further, in step **1106**, in response to successfully determining that current configuration of the server (e.g., the version of the package and/or one or more configuration settings) is to be updated, the configuration controller transmits a heartbeat response to the server supervisor. The heartbeat response message includes second metadata indicative of a new configuration of the server that is to be downloaded by the server supervisor. It is appreciated that the server supervisor upon receiving the heartbeat response message downloads an updated package from the content delivery network. Moreover, it is appreciated that in case the configuration controller determines that the current package is not required to be updated i.e., the version of the package and the one or more configuration settings of the package are up to date, then the configuration controller transmits the configuration response message to the server supervisor simply indicating that no updates are required.

[0081] FIG. **12** illustrates a simplified flowchart **1200** depicting steps performed by a server service deployed in a cloud environment according to certain implementations. Specifically, FIG. **12** illustrates a flowchart depicting the steps performed by a cloud-based enterprise service in updating one or more configuration settings of a server. In some implementations, the processing depicted in FIG. **12** may be implemented by the server configuration service **347** and the configuration controller **345** (i.e., components of the cloud-based enterprise service) depicted in FIG. **3**. The processing depicted in FIG. **12** may be implemented in software (e.g., code, instructions, program) executed by one or more processing units (e.g., processors, cores) of the respective systems, hardware, or combinations thereof. The software may be stored on a non-transitory storage medium (e.g., on a memory device). The method presented in FIG. **12** and described below is intended to be illustrative and non-limiting. Although FIG. **12** depicts the various processing steps occurring in a particular sequence or order, this is not intended to be limiting. In certain alternative implementations, the steps may be performed in some different order or some steps may also be

performed in parallel.

[0082] The process commences in step **1202**, where the server configuration service (deployed in the cloud environment) receives a request from a user interface (e.g., admin UI **307** of FIG. **3**) to apply one or more updated configuration settings to a server(s) in a customer enterprise. The process of FIG. **12** then moves to step **1204**, where the server configuration service publishes the updated one or more configuration settings to the configuration controller. Thereafter, the process moves to step **1206**, where the configuration controller transmits, in response to receiving a heartbeat message from each of the one or more server supervisors (associated with the one or more servers), a heartbeat response message including the updated configuration settings. In some implementations, the cloud service may utilize a deployment orchestrator (e.g., deployment orchestrator **349** of FIG. **3**) to update the configuration settings on the one or more servers in an incremental fashion. Specifically, the deployment orchestrator may instruct the configuration controller to update the configuration setting on a first server and verify whether a heartbeat message transmitted from the first node is okay, before proceeding to update the configuration setting on subsequent server. The deployment orchestrator may also be utilized to roll back the configuration settings of the server(s) to a prior configuration setting in case of failures encountered in the process of updating the configuration settings of the server(s).

[0083] Embodiments described above provide exemplary mechanisms for a customer enterprise deployment to be fully managed by a cloud services infrastructure. For example, FIG. **3** describes an exemplary architecture for deploying a single server (i.e., node) in the customer enterprise **301** and having a configuration controller (e.g., CBECC **345**) included in the cloud services infrastructure **340** that manages deployment of packages on the server.

[0084] According to some embodiments, certain applications may require a capability of deploying a fleet of servers (i.e., a cluster of servers) in the customer enterprise **301**. A naïve approach of creating the cluster of nodes includes utilizing a UI (e.g., UI **307** in FIG. **3**) to create the nodes in an individual fashion. In such an approach, a service principal is set up for the node by a configuration component associated with the UI. Such an approach has several drawbacks. For instance, all nodes in the cluster share the same service principal, and moreover such an approach exposes a private key (e.g., in the UI) for different entities to access, thereby incurring security issues. Embodiments discussed below overcome the above stated drawbacks and provide for a cloud services infrastructure that is configured to process a single installation command that can deploy and manage the cluster of servers. Additionally, each server in the cluster of servers is assigned a unique identifier (i.e., node ID), as well as be associated with a unique service principal. In doing so, enables the cloud services infrastructure to monitor/track and audit requests issued from each server in the cluster of servers in a seamless manner. In what follows, there is described an exemplary framework for enabling the cloud services infrastructure to deploy and manage a cluster of servers, where each server is allocated a unique identifier as well as a service principal. Additionally, the framework for deploying the cluster of servers ensures that private keys generated by a particular server in the cluster of servers is not exposed to other servers or entities in the cloud services infrastructure.

[0085] FIG. **13** depicts an exemplary swim diagram illustrating steps performed for creating a cluster of servers according to certain embodiments. Specifically, FIG. **13** illustrates interactions between a user **1301**, a configuration component **1303**, an identity access controller (IAC) **1305**, and a cloud-based enterprise configuration controller **1307** (e.g., configuration controller **345** of FIG. **3**). It is noted that user **1301** corresponds to a user of a customer (e.g., an administrator), whereas the configuration component **1303** may be a module included in configuration service of the cloud services infrastructure e.g., server configuration service **347** of FIG. **3**. Further the IAC module **1305** may correspond to an entity of the cloud services infrastructure that is responsible for identity and access management related matters.

[0086] In step **S1**, the user **1301** may utilize an UI (e.g., admin UI **307** in FIG. **3**) to issue a request

for creating a cluster of servers. The issued request is directed to the configuration component **1303**. In step **S2**, the configuration component **1303** requests the IAC module **1305** to create a role (associated with some predetermined permissions) for the cluster of servers. In step **S3**, the configuration component **1303** notifies the IAC module **1305** to assign a particular set of servers (i.e., create a group) that is to be associated with the role created in step **S2**.

[0087] In step **S4**, the configuration component **1303** transmits a request to the cloud-based enterprise configuration controller **1307**. The request corresponds to the creation of a supervisor group that is to be associated with the cluster of servers. Note that as described previously, each server includes a server supervisor that is responsible for deploying packages on the server. The request in **S4** may include an identifier (e.g., group name) for the cluster of servers, as well as a package that is to be deployed on the cluster of servers by the cloud-based enterprise configuration controller **1307**.

[0088] In step **S5**, the configuration component **1303** notifies the user **1301** regarding the creation of the cluster of servers. In some implementations, the configuration component may notify the user with a plurality of attributes associated with the cluster of servers such as a tenancy information of the cluster of servers, a supervisor group identifier for the cluster of servers, and a token e.g., a boot-token or user's token, which is a short-lived token and used for creation of servers.

[0089] FIG. **14** depicts a swim diagram illustrating steps performed in provisioning a cluster of servers according to certain embodiments. The process depicted in FIG. **14** illustrates the provisioning process with respect to a single server. It is appreciated that the process may be replicated on each server included in the cluster of servers. FIG. **14** depicts interactions between an edge provisioner **1401**, a cloud-based enterprise configuration controller **1403** (e.g., configuration controller **345** of FIG. **3**), an IAC module **1405**, a content delivery network (CDN) **1407** (e.g., CDN network **341** of FIG. **3**) and a node supervisor **1407** (e.g., server supervisor **304** of FIG. **3**). It is noted that the edge provisioner **1401** is a component that is installed in each node (i.e., server) included in the customer enterprise for purposes of performing the provisioning process as described below.

[0090] The process commences in step **S1**, where the edge provisioner **1401** loads one or more inputs including a token (e.g., boot-token associated with a user), tenancy information of the cluster of servers, a supervisor group ID. It is noted that the inputs are provided to the user in response to the user issuing a request to create the cluster of servers (e.g., in step **S5** of FIG. **13**). Upon loading the inputs, the edge provisioner **1401** in step **S2**, constructs a URL for the cloud-based enterprise configuration controller **1403** based on the tenancy information. In step **S3**, the edge provisioner **1401** generates a key pair including a private key and a public key.

[0091] In step **S4**, the edge provisioner **1401** transmits a request to the cloud-based enterprise configuration controller **1403**, where the request corresponds to provisioning a service principal for the node (i.e., server). It is appreciated that the request may include the token (e.g., boot-token associated with a user), the supervisor group ID, and the public key that is generated in step **S3**.

[0092] The process then moves to step **S5**, where the cloud-based enterprise configuration controller **1403** validates the request (received from the edge provisioner **1401**) based on verifying the supervisor group ID included in the request. Responsive to a successful validation, in step **S6**, the cloud-based enterprise configuration controller **1403** creates a service principal and registers the public key (received in the request) with the IAC module **1405**. In step **S7**, the cloud-based enterprise configuration controller **1403** receives from the IAC module, an identifier for the service principal (denoted here as a service principal ID). In step **S8**, the cloud-based enterprise configuration controller **1403** associates the service principal (created in step **S6**) with a tenancy of the cluster of servers, and further forwards the service principal ID to the edge provisioner **1401** (step **S9**).

[0093] In step **S10**, the edge provisioner **1401** generates a supervisor ID (i.e., a identifier to be

associated with a supervisor of the server, also referred to herein as a node identifier). In some implementations, the edge provisioner **1401** may generate the supervisor ID based on a combination of the service principal ID and the private key that is generated in **S3**. In this manner, a unique service principal and a unique supervisor ID is generated for each server in the cluster of servers. In step **S11**, the edge provisioner **1401** downloads a package to be installed on the node (i.e., server) from the CDN **1407** included in the cloud services infrastructure and instantiates the node supervisor **1407** (step **S12**).

[0094] In step **S13**, the cloud-based enterprise configuration controller **1403** receives a heartbeat message from the node supervisor **1407**. It is noted that as described previously, upon instantiating the node supervisor in step **S12**, the node supervisor **1407** may transmit heartbeat messages in a periodic manner (i.e., at a predetermined transmission frequency e.g., once every couple of minutes). In response to receiving the heartbeat message, the cloud-based enterprise configuration controller **1403** registers the node supervisor in step **S14**, and further transmits a package that is to be installed on the node (i.e., server) in step **S15**. In turn, in step **S16**, the node supervisor **1407** executes the package received in step **S15** from the cloud-based enterprise configuration controller **1403**. Thus, embodiments of the present disclosure provide for a means to deploy a cluster of servers in a seamless manner, where each server included in the cluster of servers is assigned a unique service principal as well as a unique identifier (e.g., supervisor ID). Moreover, it is appreciated that the framework as described above for generating the cluster of servers is secure in that a private key (associated with a particular server) is never exchanged with other servers included in the cluster of servers.

[0095] As discussed in the embodiments with reference to FIG. 3, a package in a customer enterprise implementation that is managed by a cloud-based enterprise service includes a lifecycle comprising a series of steps. For instance, in step **1**, a package author writes the code associated with the package, tests the code, and upon successful testing uploads a signed tarball (i.e., a set of files packaged together in a single file) to a content delivery network (e.g., package CDN **341** of FIG. 3). In step **2**, the package author may publish the package to a configuration controller (e.g., CBECC **345** of FIG. 3) and registers for instance, a version number, a URL path, etc. In doing so, the cloud based enterprise service is aware of the existence of the package. Further, in step **3**, a release manager uses an API associated with the configuration controller to trigger a deployment of a specific version of the package. At this point of time, the deployment orchestrator (e.g., deployment orchestrator **349** of FIG. 3) commences to deploy the package to supervisor nodes (e.g., server/node supervisor **304** of FIG. 3). In what follows, there is described in detail an architecture of the deployment orchestrator that provides for deployment of a package using different deployment strategies e.g., gradual release, parallel release, etc.

[0096] In some implementations, the act of creating a package deployment sets a desired version of the package that all supervisors (i.e., nodes) across the customer enterprise need to execute/run. This implies that the desired version of a package for any supervisor is a latest deployment e.g., latest version, of the package. For example, a particular node may have version 2.0 of a particular package that is currently installed on it, and a deployment of that package may correspond to a version 3.0 of the package that is desired to be deployed on the node.

[0097] The deployment orchestrator of the present disclosure is configured to track flow of individual packages across different nodes in the customer's enterprise. Additionally, the deployment orchestrator is configured to orchestrate new package versions to customer environments in a specific manner (i.e., strategy of deploying the package such as a gradual manner or a parallel manner) as defined by package owners, potentially taking customer maintenance windows into account. It is appreciated that customer maintenance windows correspond to time periods where the customer may prohibit deployment of any new packages (e.g., for system maintenance purposes). However, it is appreciated that in such a setting, a package may go from version v.1 to version v.5 for a given supervisor in a single step (e.g., in the case where versions

v.2-v.4 were deployed during times corresponding to the maintenance window). Additionally, as will be discussed here, the deployment orchestrator provides for a functionality of automatic rollbacks i.e., the deployment orchestrator reverts the package to a previously successfully deployed version of a package in case the desired version of the package that is being installed on the node fails.

[0098] Turning to FIG. 15, there is depicted an exemplary schematic illustrating components of a deployment orchestrator according to certain embodiments. The deployment orchestrator **1510** of FIG. 15 includes a poller **1512**, a plurality of custom resources **1514**, and a controller **1516**. In some implementations, the deployment orchestrator **1510** operates as a multi-level orchestrator i.e., the deployment orchestrator performs processing associated with deployment of a package across a plurality of hierarchical levels. For instance, the plurality of hierarchical levels includes a global level, a region level, a tenancy level, and a node level. Referring to FIG. 15, the poller **1512** of the deployment orchestrator **1510**, polls the configuration controller **1520** (e.g., CBECC **345** of FIG. 3) to identify presence of any new package deployments e.g., new versions of package(s) that may have uploaded by a user and are required to be deployed across different nodes in the enterprise system.

[0099] Further, the poller **1512** is configured to create custom resource objects **1514** (referred to herein as resource objects) corresponding to the plurality of hierarchical levels. In one implementation, the poller **1512** creates a plurality of types of resource objects, where each hierarchical level of the plurality of hierarchical levels is associated with a type of resource object included in the plurality of types of resource objects. For example, the type of resource object associated with the global level is referred to as a global resource object, the type of resource object associated with the region level is a region resource object, the type of resource object associated with the tenancy level is a tenancy resource object, and the type of resource object associated with the node level is a node resource object. It is appreciated that one or more instances of each type of resource object may correspond to each hierarchical level of the plurality of hierarchical levels. For instance, considering the tenancy level as an example, the poller creates a first instance of the tenancy resource object associated with a first customer, and a second instance of the tenancy resource object associated with a second customer. It is noted that both the first instance of the tenancy resource object and the second instance of the tenancy resource object correspond to the tenancy level.

[0100] Incorporating the above stated plurality of types of resource objects enables the deployment orchestrator **1510** to monitor/process orchestration of package deployment at different levels of granularity. For example, at the tenancy level, the deployment orchestrator **1510** can adhere to per tenancy maintenance window restrictions. As another example, region level orchestration may be required, as an end state of a customer enterprise system is likely to be distributed, and thus orchestration can be performed in each region in an independent manner. Furthermore, the plurality of hierarchical levels allow for the deployment orchestrator to monitor deployment progress in different levels of granularities and identify sources of failures in a seamless manner.

[0101] Referring back to FIG. 15, in one implementation, the controller **1516** is configured to receive the different instances of the different types of resource objects. The controller **1516** further processes each of the different instances and executes necessary actions. For example, if a desired version of a package is to be deployed at a certain node, the controller **1516** communicates with the configuration controller **1520** associated with the node and commences the process of deploying the package at the node. It is appreciated that in the above described framework of the plurality of hierarchical levels, the controller **1516** deploys packages at the node level and monitors deployment of the package for each node to obtain for instance, a metric associated with the node e.g., whether or not the package deployment was successful. The metrics obtained at the node level may be aggregated and passed over to the tenancy level. In this case, a metric associated with the tenancy level may correspond to a number of nodes in the tenancy where the package was

successfully deployed. In a similar manner, the metrics aggregated at the tenancy level may be forwarded over to the region level, and thereafter to the global level in order to obtain the statistics with respect to each node deployed in different tenancies or regions in the enterprise system.

[0102] In what follows, there is initially provided a description of the different types of resource objects followed by a system level operation of the deployment orchestrator with reference to FIG. **16**. An exemplary swim diagram illustrating process steps included in rolling out (i.e., deploying) a new package is described with reference to FIG. **17**.

[0103] According to some embodiments, the type of resource object associated with the global level is a global resource object. The global resource object may include information pertaining to (i) a desired package attribute that corresponds to the package that is to be installed in a plurality of regions e.g., the desired package attribute may correspond to a particular version of the package that is received for deployment, (ii) a status attribute corresponding to a status of package deployments on one or more nodes included in the plurality of regions. Similarly, the type of resource object associated with the region level is a region resource object. The region resource object may include information pertaining to: (i) a name of a region, (ii) a desired package attribute corresponding to the package that is to be installed in the region, and (iii) a status attribute corresponding to a status of package deployments on one or more nodes included in the region.

[0104] By some embodiments, the type of resource object associated with the tenancy level is a tenancy resource object, which may include information pertaining to: (i) a name of a tenancy, (ii) a desired package attribute corresponding to the package that is to be installed in the tenancy, (iii) a status attribute corresponding to a status of package deployments on one or more nodes included in the tenancy, and (iv) a maintenance attribute identifying a maintenance window associated with the tenancy. It is appreciated that in some implementations, a region may include one or more cells (i.e., different geographical areas within the region. In such a case, the tenancy resource object may include a cell attribute corresponding to a cell of a region that the tenancy belongs to.

[0105] By some embodiments, the type of resource object associated with the node level is a node resource object, which may include information pertaining to: (i) an identifier of a node, (ii) a desired package attribute corresponding to the package that is to be installed on the node, (iii) a current package attribute corresponding to a current package that is currently executed on the node (iv) a tenant attribute corresponding to a tenancy to which the node belongs, and (v) a status attribute corresponding to a status of package deployment in the node.

[0106] Further, as stated previously, in the deployment of a particular package across the different hierarchical levels, performance statistics may be observed at each hierarchical level. In some implementations, a success criteria and/or a failure criteria may be associated with each hierarchical level. It is appreciated that the success/failure criteria associated with a particular hierarchical level e.g., tenancy level, may be preconfigured independently (and be different) from the success/failure criteria associated with other levels e.g., region level. In some implementations, at every hierarchical level of orchestration, the failure criteria (associated with the hierarchical level) may be evaluated to make a decision on whether the package deployment should proceed for additional nodes at that level. For example, a failure criteria (e.g., number of nodes in a failed state being 20 nodes or higher) could be met for a specific tenancy, where further deployment of the package on other nodes in the tenancy is halted. However, it is noted that in doing so, would not automatically mean that the region rollout would stop for other tenancies in the region, unless the failure criteria associated with the region is met.

[0107] In some embodiments, the controller **1516** of the deployment orchestrator **1510** may include a plurality of sub-controllers. For example, the controller **1516** may include a sub-controller that is associated with each hierarchical level of the plurality of hierarchical levels. In such a setting, the sub-controller associated with a particular hierarchical level is configured to process instances of the resource object type associated with the particular hierarchical level. For example, a node level sub-controller processes all instances of the node resource object whereas a tenancy level sub-

controller processes all instances of the tenancy resource object. The sub-controllers may be utilized in a variety of ways to achieve different deployment strategies.

[0108] According to some embodiments, the deployment of a new package at a particular hierarchical level is commenced based on a time of updating instances of resource object(s) corresponding to the particular level. It is appreciated that in some implementations, the new package commences deployment upon the desired package attribute (included in the instance of the resource object) being updated to store a value corresponding to, for example, a version of the new package. Thus, the timing of deployments at different hierarchical levels may be controlled (e.g., by sub-controllers associated with each of the different hierarchical levels), by updating the desired package attribute in the different instances of the resource objects at different times. For example, considering the tenancy level, processing associated with deploying the package in a first tenancy (associated with a first instance of the tenancy resource object) may commence at a first time instance that is earlier than a second time instance corresponding to a start time of deploying the package a second tenancy (associated with a second instance of the tenancy resource object). This can be accomplished by the tenancy level sub-controller updating the desired package attribute field of the first instance of the tenancy resource object at a first time instant and updating the desired package attribute field of the second instance of the tenancy resource object at a second time instant, where the second time instant is later than the first time instant.

[0109] Furthermore, the deployment orchestrator of the present disclosure also provisions for a parallel deployment of a package e.g., deploying the package simultaneously in two tenancies by updating the respective desired package attribute fields at the same time. It is noted that although the above description with regard to achieving different deployment strategies is described with reference to sub-controllers, it is in no way limiting the scope of the present disclosure. For example, a single controller e.g., controller **1516** may be configured to perform the functions of the individual sub-controllers. Additionally, in some implementations, mechanisms other than setting of the desired package attribute field may be utilized to commence deployment of the package. For example, each of the types of resource objects may include a time attribute which corresponds to a time for commencing package deployment. The sub-controllers (or alternatively, the controller **1516**) may set different time values corresponding to the time attributes of different resource objects to achieve different package deployment timings.

[0110] FIG. **16** depicts an exemplary system diagram illustrating operations of a deployment orchestrator according to certain embodiments. An enterprise deployment environment may be spread geographically across several regions. For instance, as shown in FIG. **16**, the deployment environment includes a first enterprise deployed in region 1 **1610**, a second enterprise deployed in region 2 **1620**, and a third enterprise deployed in region K **1650**. It is noted that each region is associated with a cloud based enterprise configuration controller (i.e., CBECC) e.g., CBECC **345** of FIG. **3**, that handles deployments of packages in nodes included in the region. For instance, as shown in FIG. **16**, region 1 **1610** is associated with CBECC **1611**, region 2 **1620** is associated with CBECC **1621**, and region K **1650** is associated with CBECC **1651**. As stated previously, the CBECC of a particular region is also referred to herein as a configuration controller. Each configuration controller includes a database that stores information related to all nodes in its region.

[0111] In some implementations, in the scenario of having multiple regions as depicted in FIG. **16**, one of the regions (e.g., region 1 **1610**) is designated as a principal region. The principal region i.e., region 1, includes a deployment orchestrator **1607**. The deployment orchestrator **1607** may correspond to the deployment orchestrator **1510** of FIG. **15**. As region 1 is designated as the principal region, the deployment orchestrator included in region 1 (i.e., deployment orchestrator **1607**) is designated as a global deployment orchestrator. According to some embodiments, the global deployment orchestrator (i.e., deployment orchestrator **1607**) is responsible for managing/monitoring new package deployments to all nodes across all regions.

[0112] The deployment orchestrator **1610** includes a poller **1602** and a controller **1604**. The

functionalities of the poller **1602** and the controller **1604** correspond to the functionalities of poller **1512** and the controller **1516** (included in deployment orchestrator **1510**), respectively, as described previously with reference to FIG. **15**. Essentially, the poller **1602** is responsible for collecting information (e.g., how many nodes are included in a region, how many tenancies are present in the region, a health of the nodes in the region, etc.) from the different regions in order to build a system view with respect to the nodes across different regions i.e., a global truth of information of a system. It is appreciated that the poller **1602** may be provided (e.g., by a service) a list of regions at deployment time.

[0113] In some implementations, the poller **1602** creates the global source of truth by interrogating with all the regions and requesting any changes made in the regions (e.g., deployment of new nodes) to be delivered to it. For instance, the poller **1602** may poll the different configuration controllers (e.g., by transmitting polling messages) associated with the different regions to obtain information about the regions. The information may relate to whether new version(s) of any packages are available for deployment. In doing so, it is noted that the poller **1602** may also identify the different levels of hierarchy included in the system. As shown in FIG. **16**, poller **1602** polls the configuration controller **1621** of region 2, as well as the configuration controller **1651** of region K **1650** to obtain information (e.g., events and/or new package versions available for deployment) from the regions.

[0114] Once the poller **1602** receives such information from the different regions, the poller **1602** may create/update different resource objects (e.g., global resource object, region resource object, tenancy resource object, and node resource object) associated with different hierarchical levels as described previously with reference to FIG. **15**. In the case of a new version of package being available for deployment, the controller **1604** processes the different resource objects and rolls out the new package version for deployment (on the nodes) in the different regions. Further, the controller **1604** may also transmit heartbeat request messages to the configuration controllers of the different regions. For instance, as shown in FIG. **16**, the controller **1604** transmits heartbeat request messages to the configuration controller **1621** of region 2 as well as heartbeat request messages to configuration controller **1651** of region K. It is appreciated that the heartbeat request messages correspond to requests issued by the controller **1604** in order to obtain information regarding a status of deployment of the package(s) across different regions. In this manner, the deployment orchestrator **1610** manages and monitors deployment of a particular package(s) across different nodes included in the different regions. It is appreciated that poller **1602** and the controller **1604** of the deployment orchestrator **1607** also communicate with the configuration controller **1611** associated with region 1 **1610** in a manner similar to how they communicate with configuration controllers associated with other regions.

[0115] FIG. **17** depicts a swim diagram illustrating steps performed in deploying a package according to certain embodiments. Specifically, FIG. **17** depicts interactions between a user **1701**, a cloud based enterprise configuration controller (CBECC) **1703** (also referred to herein as a configuration controller), a poller **1705**, resource objects **1707**, a controller **1709**, and a node **1711**. It is noted that the CBECC **1703** corresponds to the configuration controller associated with a principal region (e.g., CBECC **1611** associated with region 1 in FIG. **16**). Further, the poller **1705** and the controller **1709** are components of a deployment orchestrator that is associated with the principal region e.g., deployment orchestrator **1607** of FIG. **16**.

[0116] The process commences in step S1, where user **1701** utilizes an API associated with the configuration controller **1703** to submit a new package for deployment. An entry associated with the issuing of the new package for deployment is also stored in a database associated with the configuration controller **1703**. The configuration controller **1703** in step S2 may respond to the user **1701** with an acknowledgement message.

[0117] In step S3, the poller **1705** polls the configuration controller **1703** to check for presence of new package deployments. In step S4, the configuration controller **1703** provides the poller **1705**

with a list of any new packages that are available for deployment. The process then moves to step S5, where the poller **1705** process the packages for deployments and verifies for each package whether the package is a new version that is ready for deployment. If the poller **1705** successfully verifies that a new version of the package is available, then the poller updates a global resource object (associated with the package) **1707**. It is appreciated that updating the global resource object may correspond to updating the desired package attribute to reflect a version of the new package that is to be deployed. Note that the global resource object is also referred to herein as a global customer resource object.

[0118] In step S6, the controller **1709** receives updates performed on the global resource object. In step S7, the controller processes the global resource object and updates the region resource objects i.e., instance of resource objects corresponding to the next level in the plurality of hierarchical levels. In one implementation, the region resource objects may be updated to reflect that the desired package attribute corresponds to the version of the new package that is to be deployed. It is noted that as stated previously, the controller **1709** may include multiple sub-controllers e.g., a global sub-controller, a region sub-controller, a tenancy sub-controller, and a node sub-controller that are configured to process resource objects corresponding to the global level, region level, the tenancy level, and the node level, respectively. As such, the processing depicted in step S7 may be performed by the global sub-controller.

[0119] In step S8, the region sub-controller may receive events associated with instances of region resource objects and updates (in step S9) the resource objects associated with the tenancy level (i.e., the level below the region level in the hierarchical levels). In a similar manner, in step S10, the tenancy sub-controller may receive events associated with instances of tenancy resource objects and updates (in step S11) the resource objects associated with the node level (i.e., the level below the tenancy level in the hierarchical levels). Further, in step S12, the node sub-controller receives events associated with the instances of node resource objects and processes the events (e.g., availability of new package version to be deployed) in step S13.

[0120] In the case that there is a new package version to be deployed, the controller **1709** (e.g., the node sub-controller included in the controller **1709**) issues a call to an API of the configuration controller **1703** to update package configuration to reflect the new version of the package as the desired version that is to be deployed. It is noted that the call depicted in step S14 may be a cross-regional calls made to nodes in other regions. In step S15, the controller **1709** (e.g., the node sub-controller included in the controller **1709**) receives an acknowledgment response.

[0121] Once the package configurations are updated at the node(s), the node **1711** initiates a heartbeat message to be sent to the configuration controller **1703** (step S16). In response, the configuration controller **1703** provides a heartbeat response message in step S17, where the heartbeat response message includes metrics related to the package. In step S18, the controller **1709** transmits a request to the configuration controller **1703** to receive heartbeat message(s) sent to the configuration controller **1703** by the node **1711**. This is done in order for the controller to determine whether the package was successfully deployed on the node. In step S19, the configuration controller **1703** may transmit the heartbeat messages to the controller **1709** along with an acknowledgement response.

[0122] In step S20, the controller processes the heartbeat messages to analyze metrics associated with the node (e.g., health check metrics) to determine whether the package was successfully deployed on the node. Upon successful determination, the controller **1709** may update resource object associated with the node to reflect a status of correct deployment of the package on the node. It is noted that the process as described above is a continuous process until all nodes in the system have been updated to reflect that the new version of the package is successfully deployed on the nodes. Further, in step S21, upon successful deployment of the package on all nodes, the controller issues a call to the configuration controller to update a status of deployment (with respect to the new package) as being successful, and receives an acknowledgement in step S22.

[0123] FIG. **18** depicts an exemplary state transition diagram corresponding to one hierarchical level of a plurality of hierarchical levels according to certain embodiments. Specifically, FIG. **18** depicts a state transition diagram corresponding to a node level. As shown in FIG. **18**, for a particular node, upon receiving a new package for deployment, the node enters a created state **1801** i.e., node is ready for package to be deployed. From the created state **1801**, one of two possible transitions may occur. The node may transition to a stopped state **1803** or the node may transition to a deploying state **1805**.

[0124] The node transitions into state **1803**, where it stops the deployment of the new package. Such a scenario may occur in the case where while the node is about to commence deployment of the new package, another new package for deployment is received. Thus, the node stops the deployment of the previous package and enters the stopped state **1803**. However, if no other new package is present, the node transitions from state **1801** to state **1805** (i.e., deploying state). In this case, the node initiates a call with an API of the configuration controller to initiate deployment of the new package.

[0125] From the deploying state **1805**, the node may transition to one of two states i.e., a failed state **1807** or a deployed state **1809**. In the deploying state **1805**, once the process of deploying the new package has commenced, in one implementation, the node monitors metrics for the package e.g., health check failures, process restarts, resource usage etc. As stated previously, each level of the hierarchical levels is associated with a success criteria and a failure criteria. Such success and failure criteria determine whether package deployment is to proceed. When the node is in the deploying state **1805**, if the node observes that the monitored metrics are unhealthy i.e., failure criteria associated with the node level is met, then the node transitions to the failed state and stops deployment of the new package on the node. However, when the node is in the deploying state **1805** and observes that the monitored metrics are healthy i.e., success criteria associated with the node level is met, then the node transitions to the deployed state **1809**.

[0126] In some implementations, after transitioning to either the stopped state **1803** or the deployed state **1809**, the node thereafter enters a wait state, where the node waits for the next new package deployment to occur. In some implementations, when the node transitions to the failed state **1807** due to a failure criteria associated with the node being satisfied, the node stops deployment of the new package and thereafter transitions to a rollback state, which is described next with reference to FIG. **19**.

[0127] FIG. **19** depicts an exemplary state transition diagram corresponding to a rollback operation according to certain embodiments. Upon the node stopping deployment of the package due to meeting a failure criteria, the node is in an unhealthy state **1901** i.e., one or more metrics are unhealthy. Thereafter, the node transitions to a rolling back state **1903**, where the node commences a rollback operation e.g., a process where the node commences to re-install a previous version of the package that was successfully executed on the node.

[0128] From the rolling back state **1903**, the node may transition to one of two states: a failed state **1907** or a rolled back state **1905**. In the rolling back state **1903**, once the rollback operation has commenced, the node monitors metrics for the package installation (i.e., previous version of the package). Specifically, the node monitors metrics to determine whether a success criteria or a failure criteria is met by the node. If the node observes that the monitored metrics are unhealthy i.e., failure criteria associated is met, then the node transitions to the failed state **1907** and stops the rollback operation. However, when the node is in the rolling back state **1903** and observes that the monitored metrics are healthy i.e., success criteria associated with the node level is met, then the node transitions to the rolled back state **1905**. It is appreciated that after transitioning to either the failed state **1907** or the rolled back state **1905**, the node thereafter enters a wait phase, where the node waits for the next new package deployment event to occur.

[0129] FIG. **20** depicts an exemplary state transition diagram corresponding to another hierarchical level of a plurality of hierarchical levels according to certain embodiments. Specially, FIG. **20**

depicts a state transition diagram at a tenancy level or at a region level. The description provided below is with reference to a tenancy level. However, it is appreciated that a similar state transition is relevant at the region level.

[0130] At the tenancy level, when a new package is ready for deployment in a particular tenancy, the tenancy enters a created state **2001** i.e., tenancy is ready for package to be deployed. From the created state **2001**, one of two possible transitions may occur. The tenancy may transition to a stopped state **2003** or the tenancy may transition to a deploying state **2005**.

[0131] The tenancy transitions into state **2003**, where it stops the deployment of the new package on each of the nodes included in the tenancy. Such a scenario may occur in the case where while the tenancy is about to commence deployment of the new package, another new package for deployment is received. Thus, the tenancy stops the deployment of the previous package and enters the stopped state **2003**. However, if no other new package is present, the tenancy transitions from state **2001** to state **2005** (i.e., deploying state). In this case, the tenancy initiates deployment of the new package on each of the nodes included in the tenancy and monitors rollout statistics associated with the tenancy i.e., monitors deployment metric(s) associated with all the nodes included in the tenancy.

[0132] From the deploying state **2005**, the tenancy may transition to one of three states i.e., a failed state **2007**, a deployed state **2009**, or the stopped state **2003**. In the deploying state **2005**, once the process of deploying the new package has commenced, in one implementation, the tenancy may transition to the stopped state **2003** if another new package is available for deployment. In this case, the tenancy stops the deployment process of the previous new package that it had received and enters the stopped state **2003**. Further, in the deploying state **2005**, once the process of deploying the new package has commenced, in one implementation, the tenancy monitors metrics for all nodes included in the tenancy. As stated previously, each level of the hierarchical levels is associated with a success criteria and a failure criteria. Such success and failure criteria determine whether package deployment is to proceed. When the tenancy is in the deploying state **2005** and observes that the failure criteria associated with a tenancy level is met, then the tenancy transitions to the failed state **2007** and stops deployment of the new package on all the nodes. However, when the tenancy is in the deploying state **2005** and observes that a success criteria associated with the tenancy level is met, then the tenancy transitions to the deployed state **2009**. In some implementations, after transitioning to either the stopped state **2003** the failed state **2007**, or the deployed state **2009**, the tenancy thereafter enters a wait state, where the tenancy waits for the next new package deployment to occur.

[0133] Entities of various types, such as companies, educational institutions, medical facilities, governmental departments, and private individuals, among other examples, operate computing environments for various purposes. Computing environments, which can also be referred to as information technology environments, can include inter-networked, physical hardware devices, the software executing on the hardware devices, and the users of the hardware and software. As an example, an entity such as a school can operate a Local Area Network (LAN) that includes desktop computers, laptop computers, smart phones, and tablets connected to a physical and wireless network, where users correspond to teachers and students. In this example, the physical devices may be in buildings or a campus that is controlled by the school. As another example, an entity such as a business can operate a Wide Area Network (WAN) that includes physical devices in multiple geographic locations where the offices of the business are located. In this example, the different offices can be inter-networked using a combination of public networks such as the Internet and private networks. As another example, an entity can operate a data center at a centralized location, where computing resources (such as compute, memory, and/or networking resources) are kept and maintained, and whose resources are accessible over a network to users who may be in different geographical locations. In this example, users associated with the entity that operates the data center can access the computing resources in the data center over public and/or private

networks that may not be operated and controlled by the same entity. Alternatively, or additionally, the operator of the data center may provide the computing resources to users associated with other entities, for example on a subscription basis. Such a data center operator may be referred to as a cloud services provider, and the services provided by such an entity may be described by one or more service models, such as a Software-as-a Service (SaaS) model, Infrastructure-as-a-Service (IaaS) model, or Platform-as-a-Service (PaaS), among others. In these examples, users may expect resources and/or services to be available on demand and without direct active management by the user, a resource delivery model often referred to as cloud computing.

[0134] Entities that operate computing environments need information about their computing environments. For example, an entity may need to know the operating status of the various computing resources in the entity's computing environment, so that the entity can administer the environment, including performing configuration and maintenance, performing repairs or replacements, provisioning additional resources, removing unused resources, or addressing issues that may arise during operation of the computing environment, among other examples. As another example, an entity can use information about a computing environment to identify and remediate security issues that may endanger the data, users, and/or equipment in the computing environment. As another example, an entity may be operating a computing environment for some purpose (e.g., to run an online store, to operate a bank, to manage a municipal railway, etc.) and may want information about the computing environment that can aid the entity in understanding whether the computing environment is operating efficiently and for its intended purpose.

[0135] Collection and analysis of the data from a computing environment can be performed by a data intake and query system such as is described herein. A data intake and query system can ingest, and store data obtained from the components in a computing environment, and can enable an entity to search, analyze, and visualize the data. Through these and other capabilities, the data intake and query system can enable an entity to use the data for administration of the computing environment, to detect security issues, to understand how the computing environment is performing or being used, and/or to perform other analytics.

[0136] FIG. 21 is a block diagram illustrating an example computing environment **2100** that includes a data intake and query system **2110**. The data intake and query system **2110** obtains data from a data source **2102** in the computing environment **2100** and ingests the data using an indexing system **2120**. A search system **2160** of the data intake and query system **2110** enables users to navigate the indexed data. Though drawn with separate boxes in FIG. #AA, in some implementations the indexing system **2120** and the search system **2160** can have overlapping components. A computing device **2104**, running a network access application **2106**, can communicate with the data intake and query system **2110** through a user interface system **2114** of the data intake and query system **2110**. Using the computing device **2104**, a user can perform various operations with respect to the data intake and query system **2110**, such as administration of the data intake and query system **2110**, management and generation of “knowledge objects,” (user-defined entities for enriching data, such as saved searches, event types, tags, field extractions, lookups, reports, alerts, data models, workflow actions, and fields), initiating of searches, and generation of reports, among other operations. The data intake and query system **2110** can further optionally include apps **2112** that extend the search, analytics, and/or visualization capabilities of the data intake and query system **2110**.

[0137] The data intake and query system **2110** can be implemented using program code that can be executed using a computing device. A computing device is an electronic device that has a memory for storing program code instructions and a hardware processor for executing the instructions. The computing device can further include other physical components, such as a network interface or components for input and output. The program code for the data intake and query system **2110** can be stored on a non-transitory computer-readable medium, such as a magnetic or optical storage disk or a flash or solid-state memory, from which the program code can be loaded into the memory of

the computing device for execution. “Non-transitory” means that the computer-readable medium can retain the program code while not under power, as opposed to volatile or “transitory” memory or media that requires power in order to retain data.

[0138] In various examples, the program code for the data intake and query system **2110** can be executed on a single computing device, or execution of the program code can be distributed over multiple computing devices. For example, the program code can include instructions for both indexing and search components (which may be part of the indexing system **2120** and/or the search system **2160**, respectively), which can be executed on a computing device that also provides the data source **2102**. As another example, the program code can be executed on one computing device, where execution of the program code provides both indexing and search components, while another copy of the program code executes on a second computing device that provides the data source **2102**. As another example, the program code can be configured such that, when executed, the program code implements only an indexing component or only a search component. In this example, a first instance of the program code that is executing the indexing component and a second instance of the program code that is executing the search component can be executing on the same computing device or on different computing devices.

[0139] The data source **2102** of the computing environment **2100** is a component of a computing device that produces machine data. The component can be a hardware component (e.g., a microprocessor or a network adapter, among other examples) or a software component (e.g., a part of the operating system or an application, among other examples). The component can be a virtual component, such as a virtual machine, a virtual machine monitor (also referred as a hypervisor), a container, or a container orchestrator, among other examples. Examples of computing devices that can provide the data source **2102** include personal computers (e.g., laptops, desktop computers, etc.), handheld devices (e.g., smart phones, tablet computers, etc.), servers (e.g., network servers, compute servers, storage servers, domain name servers, web servers, etc.), network infrastructure devices (e.g., routers, switches, firewalls, etc.), and “Internet of Things” devices (e.g., vehicles, home appliances, factory equipment, etc.), among other examples. Machine data is electronically generated data that is output by the component of the computing device and reflects activity of the component. Such activity can include, for example, operation status, actions performed, performance metrics, communications with other components, or communications with users, among other examples. The component can produce machine data in an automated fashion (e.g., through the ordinary course of being powered on and/or executing) and/or as a result of user interaction with the computing device (e.g., through the user's use of input/output devices or applications). The machine data can be structured, semi-structured, and/or unstructured. The machine data may be referred to as raw machine data when the data is unaltered from the format in which the data was output by the component of the computing device. Examples of machine data include operating system logs, web server logs, live application logs, network feeds, metrics, change monitoring, message queues, and archive files, among other examples.

[0140] As discussed in greater detail below, the indexing system **2120** obtains machine data from the data source **2102** and processes and stores the data. Processing and storing of data may be referred to as “ingestion” of the data. Processing of the data can include parsing the data to identify individual events, where an event is a discrete portion of machine data that can be associated with a timestamp. Processing of the data can further include generating an index of the events, where the index is a data storage structure in which the events are stored. The indexing system **2120** does not require prior knowledge of the structure of incoming data (e.g., the indexing system **2120** does not need to be provided with a schema describing the data). Additionally, the indexing system **2120** retains a copy of the data as it was received by the indexing system **2120** such that the original data is always available for searching (e.g., no data is discarded, though, in some examples, the indexing system **2120** can be configured to do so).

[0141] The search system **2160** searches the data stored by the indexing **2120** system. As discussed

in greater detail below, the search system **2160** enables users associated with the computing environment **2100** (and possibly also other users) to navigate the data, generate reports, and visualize search results in “dashboards” output using a graphical interface. Using the facilities of the search system **2160**, users can obtain insights about the data, such as retrieving events from an index, calculating metrics, searching for specific conditions within a rolling time window, identifying patterns in the data, and predicting future trends, among other examples. To achieve greater efficiency, the search system **2160** can apply map-reduce methods to parallelize searching of large volumes of data. Additionally, because the original data is available, the search system **2160** can apply a schema to the data at search time. This allows different structures to be applied to the same data, or for the structure to be modified if or when the content of the data changes. Application of a schema at search time may be referred to herein as a late-binding schema technique.

[0142] The user interface system **2114** provides mechanisms through which users associated with the computing environment **2100** (and possibly others) can interact with the data intake and query system **2110**. These interactions can include configuration, administration, and management of the indexing system **2120**, initiation and/or scheduling of queries that are to be processed by the search system **2160**, receipt or reporting of search results, and/or visualization of search results. The user interface system **2114** can include, for example, facilities to provide a command line interface or a web-based interface.

[0143] Users can access the user interface system **2114** using a computing device **2104** that communicates with data intake and query system **2110**, possibly over a network. A “user,” in the context of the implementations and examples described herein, is a digital entity that is described by a set of information in a computing environment. The set of information can include, for example, a user identifier, a username, a password, a user account, a set of authentication credentials, a token, other data, and/or a combination of the preceding. Using the digital entity that is represented by a user, a person can interact with the computing environment **2100**. For example, a person can log in as a particular user and, using the user's digital information, can access the data intake and query system **2110**. A user can be associated with one or more people, meaning that one or more people may be able to use the same user's digital information. For example, an administrative user account may be used by multiple people who have been given access to the administrative user account. Alternatively, or additionally, a user can be associated with another digital entity, such as a bot (e.g., a software program that can perform autonomous tasks). A user can also be associated with one or more entities. For example, a company can have associated with it a number of users. In this example, the company may control the users' digital information, including assignment of user identifiers, management of security credentials, control of which persons are associated with which users, and so on.

[0144] The computing device **2104** can provide a human-machine interface through which a person can have a digital presence in the computing environment **2100** in the form of a user. The computing device **2104** is an electronic device having one or more processors and a memory capable of storing instructions for execution by the one or more processors. The computing device **2104** can further include input/output (I/O) hardware and a network interface. Applications executed by the computing device **2104** can include a network access application **2106**, such as a web browser, which can use a network interface of the client computing device **2104** to communicate, over a network, with the user interface system **2114** of the data intake and query system **2110**. The user interface system **2114** can use the network access application **2106** to generate user interfaces that enable a user to interact with the data intake and query system **2110**. A web browser is one example of a network access application. A shell tool can also be used as a network access application. In some examples, the data intake and query system **2110** is an application executing on the computing device **2106**. In such examples, the network access application **2106** can access the user interface system **2114** without going over a network.

[0145] The data intake and query system **2110** can optionally include apps **2112**. An app of the data intake and query system **2110** is a collection of configurations, knowledge objects (a user-defined entity that enriches the data in the data intake and query system **2110**), views, and dashboards that may provide additional functionality, different techniques for searching the data, and/or additional insights into the data. The data intake and query system **2110** can execute multiple applications simultaneously. Example applications include an information technology service intelligence application, which can monitor and analyze the performance and behavior of the computing environment **2100**, and an enterprise security application, which can include content and searches to assist security analysts in diagnosing and acting on anomalous or malicious behavior in the computing environment **2100**.

[0146] Though FIG. **21** illustrates only one data source, in practical implementations, the computing environment **2100** contains many data sources spread across numerous computing devices. The computing devices may be controlled and operated by a single entity. For example, in an “on the premises” or “on-prem” implementation, the computing devices may physically and digitally be controlled by one entity, meaning that the computing devices are in physical locations that are owned and/or operated by the entity and are within a network domain that is controlled by the entity. In an entirely on-prem implementation of the computing environment **2100**, the data intake and query system **2110** executes on an on-prem computing device and obtains machine data from on-prem data sources. An on-prem implementation can also be referred to as an “enterprise” network, though the term “on-prem” refers primarily to physical locality of a network and who controls that location while the term “enterprise” may be used to refer to the network of a single entity. As such, an enterprise network could include cloud components.

[0147] “Cloud” or “in the cloud” refers to a network model in which an entity operates network resources (e.g., processor capacity, network capacity, storage capacity, etc.), located for example in a data center, and makes those resources available to users and/or other entities over a network. A “private cloud” is a cloud implementation where the entity provides the network resources only to its own users. A “public cloud” is a cloud implementation where an entity operates network resources in order to provide them to users that are not associated with the entity and/or to other entities. In this implementation, the provider entity can, for example, allow a subscriber entity to pay for a subscription that enables users associated with subscriber entity to access a certain amount of the provider entity's cloud resources, possibly for a limited time. A subscriber entity of cloud resources can also be referred to as a tenant of the provider entity. Users associated with the subscriber entity access the cloud resources over a network, which may include the public Internet. In contrast to an on-prem implementation, a subscriber entity does not have physical control of the computing devices that are in the cloud and has digital access to resources provided by the computing devices only to the extent that such access is enabled by the provider entity.

[0148] In some implementations, the computing environment **2100** can include on-prem and cloud-based computing resources, or only cloud-based resources. For example, an entity may have on-prem computing devices and a private cloud. In this example, the entity operates the data intake and query system **2110** and can choose to execute the data intake and query system **2110** on an on-prem computing device or in the cloud. In another example, a provider entity operates the data intake and query system **2110** in a public cloud and provides the functionality of the data intake and query system **2110** as a service, for example under a Software-as-a-Service (SaaS) model, to entities that pay for the user of the service on a subscription basis. In this example, the provider entity can provision a separate tenant (or possibly multiple tenants) in the public cloud network for each subscriber entity, where each tenant executes a separate and distinct instance of the data intake and query system **2110**. In some implementations, the entity providing the data intake and query system **2110** is itself subscribing to the cloud services of a cloud service provider. As an example, a first entity provides computing resources under a public cloud service model, a second entity subscribes to the cloud services of the first provider entity and uses the cloud computing resources

to operate the data intake and query system **2110**, and a third entity can subscribe to the services of the second provider entity in order to use the functionality of the data intake and query system **2110**. In this example, the data sources are associated with the third entity, users accessing the data intake and query system **2110** are associated with the third entity, and the analytics and insights provided by the data intake and query system **2110** are for purposes of the third entity's operations. [0149] FIG. **22** is a block diagram illustrating in greater detail an example of an indexing system **2220** of a data intake and query system, such as the data intake and query system **2110** of FIG. **21**. The indexing system **2220** of FIG. **22** uses various methods to obtain machine data from a data source **2202** and stores the data in an index **2238** of an indexer **2232**. As discussed previously, a data source is a hardware, software, physical, and/or virtual component of a computing device that produces machine data in an automated fashion and/or as a result of user interaction. Examples of data sources include files and directories; network event logs; operating system logs, operational data, and performance monitoring data; metrics; first-in, first-out queues; scripted inputs; and modular inputs, among others. The indexing system **2220** enables the data intake and query system to obtain the machine data produced by the data source **2202** and to store the data for searching and retrieval.

[0150] Users can administer the operations of the indexing system **2220** using a computing device **2204** that can access the indexing system **2220** through a user interface system **2214** of the data intake and query system. For example, the computing device **2204** can be executing a network access application **2206**, such as a web browser or a terminal, through which a user can access a monitoring console **2216** provided by the user interface system **2214**. The monitoring console **2216** can enable operations such as: identifying the data source **2202** for data ingestion; configuring the indexer **2232** to index the data from the data source **2232**; configuring a data ingestion method; configuring, deploying, and managing clusters of indexers; and viewing the topology and performance of a deployment of the data intake and query system, among other operations. The operations performed by the indexing system **2220** may be referred to as “index time” operations, which are distinct from “search time” operations that are discussed further below.

[0151] The indexer **2232**, which may be referred to herein as a data indexing component, coordinates and performs most of the index time operations. The indexer **2232** can be implemented using program code that can be executed on a computing device. The program code for the indexer **2232** can be stored on a non-transitory computer-readable medium (e.g., a magnetic, optical, or solid-state storage disk, a flash memory, or another type of non-transitory storage media), and from this medium can be loaded or copied to the memory of the computing device. One or more hardware processors of the computing device can read the program code from the memory and execute the program code in order to implement the operations of the indexer **2232**. In some implementations, the indexer **2232** executes on the computing device **2204** through which a user can access the indexing system **2220**. In some implementations, the indexer **2232** executes on a different computing device than the illustrated computing device **2204**.

[0152] The indexer **2232** may be executing on the computing device that also provides the data source **2202** or may be executing on a different computing device. In implementations wherein the indexer **2232** is on the same computing device as the data source **2202**, the data produced by the data source **2202** may be referred to as “local data.” In other implementations the data source **2202** is a component of a first computing device and the indexer **2232** executes on a second computing device that is different from the first computing device. In these implementations, the data produced by the data source **2202** may be referred to as “remote data.” In some implementations, the first computing device is “on-prem” and in some implementations the first computing device is “in the cloud.” In some implementations, the indexer **2232** executes on a computing device in the cloud and the operations of the indexer **2232** are provided as a service to entities that subscribe to the services provided by the data intake and query system.

[0153] For a given data produced by the data source **2202**, the indexing system **2220** can be

configured to use one of several methods to ingest the data into the indexer **2232**. These methods include upload **2222**, monitor **2224**, using a forwarder **2226**, or using Hypertext Transfer Protocol (HTTP **2228**) and an event collector **2230**. These and other methods for data ingestion may be referred to as “getting data in” (GDI) methods.

[0154] Using the upload **2222** method, a user can specify a file for uploading into the indexer **2232**. For example, the monitoring console **2216** can include commands or an interface through which the user can specify where the file is located (e.g., on which computing device and/or in which directory of a file system) and the name of the file. The file may be located at the data source **2202** or maybe on the computing device where the indexer **2232** is executing. Once uploading is initiated, the indexer **2232** processes the file, as discussed further below. Uploading is a manual process and occurs when instigated by a user. For automated data ingestion, the other ingestion methods are used.

[0155] The monitor **2224** method enables the indexing system **2202** to monitor the data source **2202** and continuously or periodically obtain data produced by the data source **2202** for ingestion by the indexer **2232**. For example, using the monitoring console **2216**, a user can specify a file or directory for monitoring. In this example, the indexing system **2202** can execute a monitoring process that detects whenever the file or directory is modified and causes the file or directory contents to be sent to the indexer **2232**. As another example, a user can specify a network port for monitoring. In this example, a monitoring process can capture data received at or transmitting from the network port and cause the data to be sent to the indexer **2232**. In various examples, monitoring can also be configured for data sources such as operating system event logs, performance data generated by an operating system, operating system registries, operating system directory services, and other data sources.

[0156] Monitoring is available when the data source **2202** is local to the indexer **2232** (e.g., the data source **2202** is on the computing device where the indexer **2232** is executing). Other data ingestion methods, including forwarding and the event collector **2230**, can be used for either local or remote data sources.

[0157] A forwarder **2226**, which may be referred to herein as a data forwarding component, is a software process that sends data from the data source **2202** to the indexer **2232**. The forwarder **2226** can be implemented using program code that can be executed on the computer device that provides the data source **2202**. A user launches the program code for the forwarder **2226** on the computing device that provides the data source **2202**. The user can further configure the forwarder **2226**, for example to specify a receiver for the data being forwarded (e.g., one or more indexers, another forwarder, and/or another recipient system), to enable or disable data forwarding, and to specify a file, directory, network events, operating system data, or other data to forward, among other operations.

[0158] The forwarder **2226** can provide various capabilities. For example, the forwarder **2226** can send the data unprocessed or can perform minimal processing on the data before sending the data to the indexer **2232**. Minimal processing can include, for example, adding metadata tags to the data to identify a source, source type, and/or host, among other information, dividing the data into blocks, and/or applying a timestamp to the data. In some implementations, the forwarder **2226** can break the data into individual events (event generation is discussed further below) and send the events to a receiver. Other operations that the forwarder **2226** may be configured to perform include buffering data, compressing data, and using secure protocols for sending the data, for example.

[0159] Forwarders can be configured in various topologies. For example, multiple forwarders can send data to the same indexer. As another example, a forwarder can be configured to filter and/or route events to specific receivers (e.g., different indexers), and/or discard events. As another example, a forwarder can be configured to send data to another forwarder, or to a receiver that is not an indexer or a forwarder (such as, for example, a log aggregator).

[0160] The event collector **2230** provides an alternate method for obtaining data from the data

source **2202**. The event collector **2230** enables data and application events to be sent to the indexer **2232** using HTTP **2228**. The event collector **2230** can be implemented using program code that can be executing on a computing device. The program code may be a component of the data intake and query system or can be a standalone component that can be executed independently of the data intake and query system and operates in cooperation with the data intake and query system.

[0161] To use the event collector **2230**, a user can, for example using the monitoring console **2216** or a similar interface provided by the user interface system **2214**, enable the event collector **2230** and configure an authentication token. In this context, an authentication token is a piece of digital data generated by a computing device, such as a server, which contains information to identify a particular entity, such as a user or a computing device, to the server. The token will contain identification information for the entity (e.g., an alphanumeric string that is unique to each token) and a code that authenticates the entity with the server. The token can be used, for example, by the data source **2202** as an alternative method to using a username and password for authentication.

[0162] To send data to the event collector **2230**, the data source **2202** is supplied with a token and can then send HTTP **2228** requests to the event collector **2230**. To send HTTP **2228** requests, the data source **2202** can be configured to use an HTTP client and/or to use logging libraries such as those supplied by Java, JavaScript, and .NET libraries. An HTTP client enables the data source **2202** to send data to the event collector **2230** by supplying the data, and a Uniform Resource Identifier (URI) for the event collector **2230** to the HTTP client. The HTTP client then handles establishing a connection with the event collector **2230**, transmitting a request containing the data, closing the connection, and receiving an acknowledgment if the event collector **2230** sends one. Logging libraries enable HTTP **2228** requests to the event collector **2230** to be generated directly by the data source. For example, an application can include or link a logging library, and through functionality provided by the logging library manage establishing a connection with the event collector **2230**, transmitting a request, and receiving an acknowledgement.

[0163] An HTTP **2228** request to the event collector **2230** can contain a token, a channel identifier, event metadata, and/or event data. The token authenticates the request with the event collector **2230**. The channel identifier, if available in the indexing system **2220**, enables the event collector **2230** to segregate and keep separate data from different data sources. The event metadata can include one or more key-value pairs that describe the data source **2202** or the event data included in the request. For example, the event metadata can include key-value pairs specifying a timestamp, a hostname, a source, a source type, or an index where the event data should be indexed. The event data can be a structured data object, such as a JavaScript Object Notation (JSON) object, or raw text. The structured data object can include both event data and event metadata. Additionally, one request can include event data for one or more events.

[0164] In some implementations, the event collector **2230** extracts events from HTTP **2228** requests and sends the events to the indexer **2232**. The event collector **2230** can further be configured to send events to one or more indexers. Extracting the events can include associating any metadata in a request with the event or events included in the request. In these implementations, event generation by the indexer **2232** (discussed further below) is bypassed, and the indexer **2232** moves the events directly to indexing. In some implementations, the event collector **2230** extracts event data from a request and outputs the event data to the indexer **2232**, and the indexer generates events from the event data. In some implementations, the event collector **2230** sends an acknowledgement message to the data source **2202** to indicate that the event collector **2230** has received a particular request from the data source **2202**, and/or to indicate to the data source **2202** that events in the request have been added to an index.

[0165] The indexer **2232** ingests incoming data and transforms the data into searchable knowledge in the form of events. In the data intake and query system, an event is a single piece of data that represents activity of the component represented in FIG. 22 by the data source **2202**. An event can be, for example, a single record in a log file that records a single action performed by the

component (e.g., a user login, a disk read, transmission of a network packet, etc.). An event includes one or more fields that together describe the action captured by the event, where a field is a key-value pair (also referred to as a name-value pair). In some cases, an event includes both the key and the value, and in some cases the event includes only the value, and the key can be inferred or assumed.

[0166] Transformation of data into events can include event generation and event indexing. Event generation includes identifying each discrete piece of data that represents one event and associating each event with a timestamp and possibly other information (which may be referred to herein as metadata). Event indexing includes storing of each event in the data structure of an index. As an example, the indexer **2232** can include a parsing module **2234** and an indexing module **2236** for generating and storing the events. The parsing module **2234** and indexing module **2236** can be modular and pipelined, such that one component can be operating on a first set of data while the second component is simultaneously operating on a second set of data. Additionally, the indexer **2232** may at any time have multiple instances of the parsing module **2234** and indexing module **2236**, with each set of instances configured to simultaneously operate on data from the same data source or from different data sources. The parsing module **2234** and indexing module **2236** are illustrated in FIG. 22 to facilitate discussion, with the understanding that implementations with other components are possible to achieve the same functionality.

[0167] The parsing module **2234** determines information about incoming event data, where the information can be used to identify events within the event data. For example, the parsing module **2234** can associate a source type with the event data. A source type identifies the data source **2202** and describes a possible data structure of event data produced by the data source **2202**. For example, the source type can indicate which fields to expect in events generated at the data source **2202** and the keys for the values in the fields, and possibly other information such as sizes of fields, an order of the fields, a field separator, and so on. The source type of the data source **2202** can be specified when the data source **2202** is configured as a source of event data. Alternatively, the parsing module **2234** can determine the source type from the event data, for example from an event field in the event data or using machine learning techniques applied to the event data.

[0168] Other information that the parsing module **2234** can determine includes timestamps. In some cases, an event includes a timestamp as a field, and the timestamp indicates a point in time when the action represented by the event occurred or was recorded by the data source **2202** as event data. In these cases, the parsing module **2234** may be able to determine from the source type associated with the event data that the timestamps can be extracted from the events themselves. In some cases, an event does not include a timestamp and the parsing module **2234** determines a timestamp for the event, for example from a name associated with the event data from the data source **2202** (e.g., a file name when the event data is in the form of a file) or a time associated with the event data (e.g., a file modification time). As another example, when the parsing module **2234** is not able to determine a timestamp from the event data, the parsing module **2234** may use the time at which it is indexing the event data. As another example, the parsing module **2234** can use a user-configured rule to determine the timestamps to associate with events.

[0169] The parsing module **2234** can further determine event boundaries. In some cases, a single line (e.g., a sequence of characters ending with a line termination) in event data represents one event while in other cases, a single line represents multiple events. In yet other cases, one event may span multiple lines within the event data. The parsing module **2234** may be able to determine event boundaries from the source type associated with the event data, for example from a data structure indicated by the source type. In some implementations, a user can configure rules the parsing module **2234** can use to identify event boundaries.

[0170] The parsing module **2234** can further extract data from events and possibly also perform transformations on the events. For example, the parsing module **2234** can extract a set of fields (key-value pairs) for each event, such as a host or hostname, source or source name, and/or source

type. The parsing module **2234** may extract certain fields by default or based on a user configuration. Alternatively, or additionally, the parsing module **2234** may add fields to events, such as a source type or a user-configured field. As another example of a transformation, the parsing module **2234** can anonymize fields in events to mask sensitive information, such as social security numbers or account numbers. Anonymizing fields can include changing or replacing values of specific fields. The parsing component **2234** can further perform user-configured transformations.

[0171] The parsing module **2234** outputs the results of processing incoming event data to the indexing module **2236**, which performs event segmentation and builds index data structures.

[0172] Event segmentation identifies searchable segments, which may alternatively be referred to as searchable terms or keywords, which can be used by the search system of the data intake and query system to search the event data. A searchable segment may be a part of a field in an event or an entire field. The indexer **2232** can be configured to identify searchable segments that are parts of fields, searchable segments that are entire fields, or both. The parsing module **2234** organizes the searchable segments into a lexicon or dictionary for the event data, with the lexicon including each searchable segment (e.g., the field “src=10.10.1.1”) and a reference to the location of each occurrence of the searchable segment within the event data (e.g., the location within the event data of each occurrence of “src=10.10.1.1”). As discussed further below, the search system can use the lexicon, which is stored in an index file **2246**, to find event data that matches a search query. In some implementations, segmentation can alternatively be performed by the forwarder **2226**.

Segmentation can also be disabled, in which case the indexer **2232** will not build a lexicon for the event data. When segmentation is disabled, the search system searches the event data directly.

[0173] Building index data structures generates the index **2238**. The index **2238** is a storage data structure on a storage device (e.g., a disk drive or other physical device for storing digital data). The storage device may be a component of the computing device on which the indexer **2232** is operating (referred to herein as local storage) or may be a component of a different computing device (referred to herein as remote storage) that the indexer **2238** has access to over a network. The indexer **2232** can manage more than one index and can manage indexes of different types. For example, the indexer **2232** can manage event indexes, which impose minimal structure on stored data and can accommodate any type of data. As another example, the indexer **2232** can manage metrics indexes, which use a highly structured format to handle the higher volume and lower latency demands associated with metrics data.

[0174] The indexing module **2236** organizes files in the index **2238** in directories referred to as buckets. The files in a bucket **2244** can include raw data files, index files, and possibly also other metadata files. As used herein, “raw data” means data as when the data was produced by the data source **2202**, without alteration to the format or content. As noted previously, the parsing component **2234** may add fields to event data and/or perform transformations on fields in the event data. Event data that has been altered in this way is referred to herein as enriched data. A raw data file **2248** can include enriched data, in addition to or instead of raw data. The raw data file **2248** may be compressed to reduce disk usage. An index file **2246**, which may also be referred to herein as a “time-series index” or tsidx file, contains metadata that the indexer **2232** can use to search a corresponding raw data file **2248**. As noted above, the metadata in the index file **2246** includes a lexicon of the event data, which associates each unique keyword in the event data with a reference to the location of event data within the raw data file **2248**. The keyword data in the index file **2246** may also be referred to as an inverted index. In various implementations, the data intake and query system can use index files for other purposes, such as to store data summarizations that can be used to accelerate searches.

[0175] A bucket **2244** includes event data for a particular range of time. The indexing module **2236** arranges buckets in the index **2238** according to the age of the buckets, such that buckets for more recent ranges of time are stored in short-term storage **2240** and buckets for less recent ranges of

time are stored in long-term storage **2242**. Short-term storage **2240** may be faster to access while long-term storage **2242** may be slower to access. Buckets may be moved from short-term storage **2240** to long-term storage **2242** according to a configurable data retention policy, which can indicate at what point in time a bucket is old enough to be moved.

[0176] A bucket's location in short-term storage **2240** or long-term storage **2242** can also be indicated by the bucket's status. As an example, a bucket's status can be “hot,” “warm,” “cold,” “frozen,” or “thawed.” In this example, hot bucket is one to which the indexer **2232** is writing data and the bucket becomes a warm bucket when the index **2232** stops writing data to it. In this example, both hot and warm buckets reside in short-term storage **2240**. Continuing this example, when a warm bucket is moved to long-term storage **2242**, the bucket becomes a cold bucket. A cold bucket can become a frozen bucket after a period of time, at which point the bucket may be deleted or archived. An archived bucket cannot be searched. When an archived bucket is retrieved for searching, the bucket becomes thawed and can then be searched.

[0177] The indexing system **2220** can include more than one indexer, where a group of indexers is referred to as an index cluster. The indexers in an index cluster may also be referred to as peer nodes. In an index cluster, the indexers are configured to replicate each other's data by copying buckets from one indexer to another. The number of copies of a bucket can be configured (e.g., three copies of each buckets must exist within the cluster), and indexers to which buckets are copied may be selected to optimize distribution of data across the cluster.

[0178] A user can view the performance of the indexing system **2220** through the monitoring console **2216** provided by the user interface system **2214**. Using the monitoring console **2216**, the user can configure and monitor an index cluster, and see information such as disk usage by an index, volume usage by an indexer, index and volume size over time, data age, statistics for bucket types, and bucket settings, among other information.

[0179] FIG. **23** is a block diagram illustrating in greater detail an example of the search system **2360** of a data intake and query system, such as the data intake and query system **2110** of FIG. **#AA**. The search system **2360** of FIG. **23** issues a query **2366** to a search head **2362**, which sends the query **2366** to a search peer **2364**. Using a map process **2370**, the search peer **2364** searches the appropriate index **2338** for events identified by the query **2366** and sends events **2378** so identified back to the search head **2362**. Using a reduce process **2382**, the search head **2362** processes the events **2378** and produces results **2368** to respond to the query **2366**. The results **2368** can provide useful insights about the data stored in the index **2338**. These insights can aid in the administration of information technology systems, in security analysis of information technology systems, and/or in analysis of the development environment provided by information technology systems.

[0180] The query **2366** that initiates a search is produced by a search and reporting app **2316** that is available through the user interface system **2314** of the data intake and query system. Using a network access application **2306** executing on a computing device **2304**, a user can input the query **2366** into a search field provided by the search and reporting app **2316**. Alternatively, or additionally, the search and reporting app **2316** can include pre-configured queries or stored queries that can be activated by the user. In some cases, the search and reporting app **2316** initiates the query **2366** when the user enters the query **2366**. In these cases, the query **2366** maybe referred to as an “ad-hoc” query. In some cases, the search and reporting app **2316** initiates the query **2366** based on a schedule. For example, the search and reporting app **2316** can be configured to execute the query **2366** once per hour, once per day, at a specific time, on a specific date, or at some other time that can be specified by a date, time, and/or frequency. These types of queries maybe referred to as scheduled queries.

[0181] The query **2366** is specified using a search processing language. The search processing language includes commands or search terms that the search peer **2364** will use to identify events to return in the search results **2368**. The search processing language can further include commands for filtering events, extracting more information from events, evaluating fields in events,

aggregating events, calculating statistics over events, organizing the results, and/or generating charts, graphs, or other visualizations, among other examples. Some search commands may have functions and arguments associated with them, which can, for example, specify how the commands operate on results and which fields to act upon. The search processing language may further include constructs that enable the query **2366** to include sequential commands, where a subsequent command may operate on the results of a prior command. As an example, sequential commands may be separated in the query **2366** by a vertical line (“|” or “pipe”) symbol.

[0182] In addition to one or more search commands, the query **2366** includes a time indicator. The time indicator limits searching to events that have timestamps described by the indicator. For example, the time indicator can indicate a specific point in time (e.g., 10:00:00 am today), in which case only events that have the point in time for their timestamp will be searched. As another example, the time indicator can indicate a range of time (e.g., the last 24 hours), in which case only events whose timestamps fall within the range of time will be searched. The time indicator can alternatively indicate all of time, in which case all events will be searched.

[0183] Processing of the search query **2366** occurs in two broad phases: a map phase **2350** and a reduce phase **2352**. The map phase **2350** takes place across one or more search peers. In the map phase **2350**, the search peers locate event data that matches the search terms in the search query **2366** and sorts the event data into field-value pairs. When the map phase **2350** is complete, the search peers send events that they have found to one or more search heads for the reduce phase **2352**. During the reduce phase **2352**, the search heads process the events through commands in the search query **2366** and aggregate the events to produce the final search results **2368**.

[0184] A search head, such as the search head **2362** illustrated in FIG. 23, is a component of the search system **2360** that manages searches. The search head **2362**, which may also be referred to herein as a search management component, can be implemented using program code that can be executed on a computing device. The program code for the search head **2362** can be stored on a non-transitory computer-readable medium and from this medium can be loaded or copied to the memory of a computing device. One or more hardware processors of the computing device can read the program code from the memory and execute the program code in order to implement the operations of the search head **2362**.

[0185] Upon receiving the search query **2366**, the search head **2362** directs the query **2366** to one or more search peers, such as the search peer **2364** illustrated in FIG. 23. “Search peer” is an alternate name for “indexer” and a search peer may be largely similar to the indexer described previously. The search peer **2364** may be referred to as a “peer node” when the search peer **2364** is part of an indexer cluster. The search peer **2364**, which may also be referred to as a search execution component, can be implemented using program code that can be executed on a computing device. In some implementations, one set of program code implements both the search head **2362** and the search peer **2364** such that the search head **2362** and the search peer **2364** form one component. In some implementations, the search head **2362** is an independent piece of code that performs searching and no indexing functionality. In these implementations, the search head **2362** may be referred to as a dedicated search head.

[0186] The search head **2362** may consider multiple criteria when determining whether to send the query **2366** to the particular search peer **2364**. For example, the search system **2360** may be configured to include multiple search peers that each have duplicative copies of at least some of the event data and are implanted using different hardware resources *q*. In this example, the sending the search query **2366** to more than one search peer allows the search system **2360** to distribute the search workload across different hardware resources. As another example, search system **2360** may include different search peers for different purposes (e.g., one has an index storing a first type of data or from a first data source while a second has an index storing a second type of data or from a second data source). In this example, the search query **2366** may specify which indexes to search, and the search head **2362** will send the query **2366** to the search peers that have those indexes.

[0187] To identify events **2378** to send back to the search head **2362**, the search peer **2364** performs a map process **2370** to obtain event data **2374** from the index **2338** that is maintained by the search peer **2364**. During a first phase of the map process **2370**, the search peer **2364** identifies buckets that have events that are described by the time indicator in the search query **2366**. As noted above, a bucket contains events whose timestamps fall within a particular range of time. For each bucket **2344** whose events can be described by the time indicator, during a second phase of the map process **2370**, the search peer **2364** performs a keyword search **2374** using search terms specified in the search query **2366**. The search terms can be one or more of keywords, phrases, fields, Boolean expressions, and/or comparison expressions that in combination describe events being searched for. When segmentation is enabled at index time, the search peer **2364** performs the keyword search **2372** on the bucket's index file **2346**. As noted previously, the index file **2346** includes a lexicon of the searchable terms in the events stored in the bucket's raw data **2348** file. The keyword search **2372** searches the lexicon for searchable terms that correspond to one or more of the search terms in the query **2366**. As also noted above, the lexicon includes, for each searchable term, a reference to each location in the raw data **2348** file where the searchable term can be found. Thus, when the keyword search identifies a searchable term in the index file **2346** that matches a search term in the query **2366**, the search peer **2364** can use the location references to extract from the raw data **2348** file the event data **2374** for each event that include the searchable term.

[0188] In cases where segmentation was disabled at index time, the search peer **2364** performs the keyword search **2372** directly on the raw data **2348** file. To search the raw data **2348**, the search peer **2364** may identify searchable segments in events in a similar manner as when the data was indexed. Thus, depending on how the search peer **2364** is configured, the search peer **2364** may look at event fields and/or parts of event fields to determine whether an event matches the query **2366**. Any matching events can be added to the event data **2374** read from the raw data **2348** file. The search peer **2364** can further be configured to enable segmentation at search time, so that searching of the index **2338** causes the search peer **2364** to build a lexicon in the index file **2346**.

[0189] The event data **2374** obtained from the raw data **2348** file includes the full text of each event found by the keyword search **2372**. During a third phase of the map process **2370**, the search peer **2364** performs event processing **2376** on the event data **2374**, with the steps performed being determined by the configuration of the search peer **2364** and/or commands in the search query **2366**. For example, the search peer **2364** can be configured to perform field discovery and field extraction. Field discovery is a process by which the search peer **2364** identifies and extracts key-value pairs from the events in the event data **2374**. The search peer **2364** can, for example, be configured to automatically extract the first **100** fields (or another number of fields) in the event data **2374** that can be identified as key-value pairs. As another example, the search peer **2364** can extract any fields explicitly mentioned in the search query **2366**. The search peer **2364** can, alternatively or additionally, be configured with particular field extractions to perform.

[0190] Other examples of steps that can be performed during event processing **2376** include: field aliasing (assigning an alternate name to a field); addition of fields from lookups (adding fields from an external source to events based on existing field values in the events); associating event types with events; source type renaming (changing the name of the source type associated with particular events); and tagging (adding one or more strings of text, or a "tags" to particular events), among other examples.

[0191] The search peer **2364** sends processed events **2378** to the search head **2362**, which performs a reduce process **2380**. The reduce process **2380** potentially receives events from multiple search peers and performs various results processing **2382** steps on the received events. The results processing **2382** steps can include, for example, aggregating the events received from different search peers into a single set of events, de-duplicating and aggregating fields discovered by different search peers, counting the number of events found, and sorting the events by timestamp (e.g., newest first or oldest first), among other examples. Results processing **2382** can further

include applying commands from the search query **2366** to the events. The query **2366** can include, for example, commands for evaluating and/or manipulating fields (e.g., to generate new fields from existing fields or parse fields that have more than one value). As another example, the query **2366** can include commands for calculating statistics over the events, such as counts of the occurrences of fields, or sums, averages, ranges, and so on, of field values. As another example, the query **2366** can include commands for generating statistical values for purposes of generating charts or graphs of the events.

[0192] The reduce process **2380** outputs the events found by the search query **2366**, as well as information about the events. The search head **2362** transmits the events and the information about the events as search results **2368**, which are received by the search and reporting app **2316**. The search and reporting app **2316** can generate visual interfaces for viewing the search results **2368**. The search and reporting app **2316** can, for example, output visual interfaces for the network access application **2306** running on a computing device **2304** to generate.

[0193] The visual interfaces can include various visualizations of the search results **2368**, such as tables, line or area charts, Choropleth maps, or single values. The search and reporting app **2316** can organize the visualizations into a dashboard, where the dashboard includes a panel for each visualization. A dashboard can thus include, for example, a panel listing the raw event data for the events in the search results **2368**, a panel listing fields extracted at index time and/or found through field discovery along with statistics for those fields, and/or a timeline chart indicating how many events occurred at specific points in time (as indicated by the timestamps associated with each event). In various implementations, the search and reporting app **2316** can provide one or more default dashboards. Alternatively, or additionally, the search and reporting app **2316** can include functionality that enables a user to configure custom dashboards.

[0194] The search and reporting app **2316** can also enable further investigation into the events in the search results **2316**. The process of further investigation may be referred to as drilldown. For example, a visualization in a dashboard can include interactive elements, which, when selected, provide options for finding out more about the data being displayed by the interactive elements. To find out more, an interactive element can, for example, generate a new search that includes some of the data being displayed by the interactive element, and thus may be more focused than the initial search query **2366**. As another example, an interactive element can launch a different dashboard whose panels include more detailed information about the data that is displayed by the interactive element. Other examples of actions that can be performed by interactive elements in a dashboard include opening a link, playing an audio or video file, or launching another application, among other examples.

[0195] FIG. **24** illustrates an example of a self-managed network **2400** that includes a data intake and query system. “Self-managed” in this instance means that the entity that is operating the self-managed network **2400** configures, administers, maintains, and/or operates the data intake and query system using its own compute resources and people. Further, the self-managed network **2400** of this example is part of the entity's on-premise network and comprises a set of compute, memory, and networking resources that are located, for example, within the confines of an entity's data center. These resources can include software and hardware resources. The entity can, for example, be a company or enterprise, a school, government entity, or other entity. Since the self-managed network **2400** is located within the customer's on-prem environment, such as in the entity's data center, the operation and management of the self-managed network **2400**, including of the resources in the self-managed network **2400**, is under the control of the entity. For example, administrative personnel of the entity have complete access to and control over the configuration, management, and security of the self-managed network **2400** and its resources.

[0196] The self-managed network **2400** can execute one or more instances of the data intake and query system. An instance of the data intake and query system may be executed by one or more computing devices that are part of the self-managed network **2400**. A data intake and query system

instance can comprise an indexing system and a search system, where the indexing system includes one or more indexers **2420** and the search system includes one or more search heads **2460**.

[0197] As depicted in FIG. **24**, the self-managed network **2400** can include one or more data sources **2402**. Data received from these data sources may be processed by an instance of the data intake and query system within self-managed network **2400**. The data sources **2402** and the data intake and query system instance can be communicatively coupled to each other via a private network **2410**.

[0198] Users associated with the entity can interact with and avail themselves of the functions performed by a data intake and query system instance using computing devices. As depicted in FIG. **24**, a computing device **2404** can execute a network access application **2406** (e.g., a web browser), that can communicate with the data intake and query system instance and with data sources **2402** via the private network **2410**. Using the computing device **2404**, a user can perform various operations with respect to the data intake and query system, such as management and administration of the data intake and query system, generation of knowledge objects, and other functions. Results generated from processing performed by the data intake and query system instance may be communicated to the computing device **2404** and output to the user via an output system (e.g., a screen) of the computing device **2404**.

[0199] The self-managed network **2400** can also be connected to other networks that are outside the entity's on-premise environment/network, such as networks outside the entity's data center. Connectivity to these other external networks is controlled and regulated through one or more layers of security provided by the self-managed network **2400**. One or more of these security layers can be implemented using firewalls **2412**. The firewalls **2412** form a layer of security around the self-managed network **2400** and regulate the transmission of traffic from the self-managed network **2400** to the other networks and from these other networks to the self-managed network **2400**.

[0200] Networks external to the self-managed network can include various types of networks including public networks **2490**, other private networks, and/or cloud networks provided by one or more cloud service providers. An example of a public network **2490** is the Internet. In the example depicted in FIG. **24**, the self-managed network **2400** is connected to a service provider network **2492** provided by a cloud service provider via the public network **2490**.

[0201] In some implementations, resources provided by a cloud service provider may be used to facilitate the configuration and management of resources within the self-managed network **2400**. For example, configuration and management of a data intake and query system instance in the self-managed network **2400** may be facilitated by a software management system **2494** operating in the service provider network **2492**. There are various ways in which the software management system **2494** can facilitate the configuration and management of a data intake and query system instance within the self-managed network **2400**. As one example, the software management system **2494** may facilitate the download of software including software updates for the data intake and query system. In this example, the software management system **2494** may store information indicative of the versions of the various data intake and query system instances present in the self-managed network **2400**. When a software patch or upgrade is available for an instance, the software management system **2494** may inform the self-managed network **2400** of the patch or upgrade. This can be done via messages communicated from the software management system **2494** to the self-managed network **2400**.

[0202] The software management system **2494** may also provide simplified ways for the patches and/or upgrades to be downloaded and applied to the self-managed network **2400**. For example, a message communicated from the software management system **2494** to the self-managed network **2400** regarding a software upgrade may include a Uniform Resource Identifier (URI) that can be used by a system administrator of the self-managed network **2400** to download the upgrade to the self-managed network **2400**. In this manner, management resources provided by a cloud service provider using the service provider network **2492** and which are located outside the self-managed

network **2400** can be used to facilitate the configuration and management of one or more resources within the entity's on-prem environment. In some implementations, the download of the upgrades and patches may be automated, whereby the software management system **2494** is authorized to, upon determining that a patch is applicable to a data intake and query system instance inside the self-managed network **2400**, automatically communicate the upgrade or patch to self-managed network **2400** and cause it to be installed within self-managed network **2400**.

[0203] Various examples and possible implementations have been described above, which recite certain features and/or functions. Although these examples and implementations have been described in language specific to structural features and/or functions, it is understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or functions described above. Rather, the specific features and functions described above are disclosed as examples of implementing the claims, and other equivalent features and acts are intended to be within the scope of the claims. Further, any or all of the features and functions described above can be combined with each other, except to the extent it may be otherwise stated above or to the extent that any such embodiments may be incompatible by virtue of their function or structure, as will be apparent to persons of ordinary skill in the art. Unless contrary to physical possibility, it is envisioned that (i) the methods/steps described herein may be performed in any sequence and/or in any combination, and (ii) the components of respective embodiments may be combined in any manner.

[0204] Processing of the various components of systems illustrated herein can be distributed across multiple machines, networks, and other computing resources. Two or more components of a system can be combined into fewer components. Various components of the illustrated systems can be implemented in one or more virtual machines or an isolated execution environment, rather than in dedicated computer hardware systems and/or computing devices. Likewise, the data repositories shown can represent physical and/or logical data storage, including, e.g., storage area networks or other distributed storage systems. Moreover, in some embodiments the connections between the components shown represent possible paths of data flow, rather than actual connections between hardware. While some examples of possible connections are shown, any of the subset of the components shown can communicate with any other subset of components in various implementations.

[0205] Examples have been described with reference to flow chart illustrations and/or block diagrams of methods, apparatus (systems), and computer program products. Each block of the flow chart illustrations and/or block diagrams, and combinations of blocks in the flow chart illustrations and/or block diagrams, may be implemented by computer program instructions. Such instructions may be provided to a processor of a general purpose computer, special purpose computer, specially-equipped computer (e.g., comprising a high-performance database server, a graphics subsystem, etc.) or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor(s) of the computer or other programmable data processing apparatus, create means for implementing the acts specified in the flow chart and/or block diagram block or blocks. These computer program instructions may also be stored in a non-transitory computer-readable memory that can direct a computer or other programmable data processing apparatus to operate in a particular manner, such that the instructions stored in the computer-readable memory produce an article of manufacture including instruction means which implement the acts specified in the flow chart and/or block diagram block or blocks. The computer program instructions may also be loaded to a computing device or other programmable data processing apparatus to cause operations to be performed on the computing device or other programmable apparatus to produce a computer implemented process such that the instructions which execute on the computing device or other programmable apparatus provide steps for implementing the acts specified in the flow chart and/or block diagram block or blocks.

[0206] In some embodiments, certain operations, acts, events, or functions of any of the algorithms

described herein can be performed in a different sequence, can be added, merged, or left out altogether (e.g., not all are necessary for the practice of the algorithms). In certain embodiments, operations, acts, functions, or events can be performed concurrently, e.g., through multi-threaded processing, interrupt processing, or multiple processors or processor cores or on other parallel architectures, rather than sequentially.

Claims

1. A method comprising: receiving, by a cloud-based configuration controller included in a cloud service infrastructure, a heartbeat message from a node deployed in a customer enterprise, the heartbeat message including first metadata indicative of a current configuration of a package that is executed by the node; determining, by the cloud-based configuration controller included in the cloud service infrastructure, whether the current configuration of the package is to be updated; responsive to successfully determining that the current configuration of the package is to be updated, transmitting, by the cloud-based configuration controller, a heartbeat response message to the node, wherein the heartbeat response message includes second metadata indicative of a new configuration of the package that is to be downloaded by the node from the cloud service infrastructure.
2. The method of claim 1, wherein the package includes an application and one or more configuration settings associated with the application.
3. The method of claim 1, wherein the first metadata includes information identifying a version of an application included in the package or values of one or more configuration settings associated with the application.
4. The method of claim 1, wherein the first metadata further includes information corresponding to a health metric of the node.
5. The method of claim 1, wherein the second metadata includes a URL associated with a content delivery network included in the cloud service infrastructure from where the node downloads the new configuration of the package.
6. The method of claim 1, wherein the new configuration of the package corresponds to a new version of an application included in the package or new values of one or more configuration settings associated with the application.
7. The method of claim 1, wherein the heartbeat message is transmitted periodically from the node to the cloud-based configuration controller included in a cloud service at a predetermined transmission frequency.
8. The method of claim 1, further comprising: receiving, by a package service deployed in the cloud service infrastructure, a request from an user interface to apply an updated configuration settings to the package executed on the node; publishing, by the package service, the updated configuration settings to the cloud-based configuration controller included in the cloud service infrastructure; and sending, by the cloud-based configuration controller included in the cloud service infrastructure, in response to receiving the heartbeat message from the node, a new heartbeat response message including the updated configuration settings.
9. The method of claim 1, further comprising: obtaining, by the cloud-based configuration controller included in the cloud service infrastructure, a first request from the node, the first request requesting creating a service principal that is to be associated with the node; generating, by the cloud-based configuration controller, a service principal identifier that is to be associated with the node; and transmitting, by the cloud-based configuration controller, the service principal identifier to the node.
10. One or more computer readable non-transitory media storing computer-executable instructions that, when executed by one or more processors, cause: receiving, by a cloud-based configuration controller included in a cloud service infrastructure, a heartbeat message from a node deployed in a

customer enterprise, the heartbeat message including first metadata indicative of a current configuration of a package that is executed by the node; determining, by the cloud-based configuration controller included in the cloud service infrastructure, whether the current configuration of the package is to be updated; responsive to successfully determining that the current configuration of the package is to be updated, transmitting, by the cloud-based configuration controller, a heartbeat response message to the node, wherein the heartbeat response message includes second metadata indicative of a new configuration of the package that is to be downloaded by the node from the cloud service infrastructure.

11. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, wherein the package includes an application and one or more configuration settings associated with the application.

12. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, wherein the first metadata includes information identifying a version of an application included in the package or values of one or more configuration settings associated with the application.

13. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, wherein the first metadata further includes information corresponding to a health metric of the node.

14. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, wherein the second metadata includes a URL associated with a content delivery network included in the cloud service infrastructure from where the node downloads the new configuration of the package.

15. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, wherein the new configuration of the package corresponds to a new version of an application included in the package or new values of one or more configuration settings associated with the application.

16. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, wherein the heartbeat message is transmitted periodically from the node to the cloud-based configuration controller included in a cloud service at a predetermined transmission frequency.

17. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, further comprising: receiving, by a package service deployed in the cloud service infrastructure, a request from an user interface to apply an updated configuration settings to the package executed on the node; publishing, by the package service, the updated configuration settings to the cloud-based configuration controller included in the cloud service infrastructure; and sending, by the cloud-based configuration controller included in the cloud service infrastructure, in response to receiving the heartbeat message from the node, a new heartbeat response message including the updated configuration settings.

18. The one or more computer readable non-transitory media storing computer-executable instructions of claim 10, further comprising: obtaining, by the cloud-based configuration controller included in the cloud service infrastructure, a first request from the node, the first request requesting creating a service principal that is to be associated with the node; generating, by the cloud-based configuration controller, a service principal identifier that is to be associated with the node; and transmitting, by the cloud-based configuration controller, the service principal identifier to the node.

19. A controller included in a cloud service infrastructure, the controller comprising: memory; and one or more processing devices coupled to the memory and configured to: receive a heartbeat message from a node deployed in a customer enterprise, the heartbeat message including first metadata indicative of a current configuration of a package that is executed by the node; determine whether the current configuration of the package is to be updated; and responsive to successfully

determining that the current configuration of the package is to be updated, transmit a heartbeat response message to the node, wherein the heartbeat response message includes second metadata indicative of a new configuration of the package that is to be downloaded by the node from the cloud service infrastructure.

20. The controller included in the cloud service infrastructure of claim 19, wherein the package includes an application, and one or more configuration settings associated with the application, and wherein the first metadata includes information identifying a version of the application included in the package or values of one or more configuration settings associated with the application.
