

# US Patent & Trademark Office

## Patent Public Search | Text View

United States Patent Application Publication

20250265081

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Sade; Raanan et al.

### SYSTEMS AND METHODS FOR COMPUTING DOT PRODUCTS OF NIBBLES IN TWO TILE OPERANDS

#### Abstract

Disclosed embodiments relate to computing dot products of nibbles in tile operands. In one example, a processor includes decode circuitry to decode a tile dot product instruction having fields for an opcode, a destination identifier to identify a M by N destination matrix, a first source identifier to identify a M by K first source matrix, and a second source identifier to identify a K by N second source matrix, each of the matrices containing doubleword elements, and execution circuitry to execute the decoded instruction to perform a flow K times for each element (M,N) of the identified destination matrix to generate eight products by multiplying each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix, and to accumulate and saturate the eight products with previous contents of the doubleword element (M,N).

**Inventors:** Sade; Raanan (Portland, OR), Rubanovich; Simon (Haifa, IL), Gradstein; Amit (Binyamina, IL), Sperber; Zeev (Zichron Yackov, IL), Heinecke; Alexander (San Jose, CA), Valentine; Robert (Kiryat Tivon, IL), Charney; Mark J. (Lexington, MA), Toll; Bret (Hillsboro, OR), Corbal; Jesus (King City, OR), Ould-Ahmed-Vall; Elmoustapha (Gilbert, AZ)

**Applicant:** Intel Corporation (Santa Clara, CA)

**Family ID:** 1000008578145

**Appl. No.:** 19/005958

**Filed:** December 30, 2024

#### Related U.S. Application Data

parent US continuation 18449651 20230814 parent-grant-document US 12182568 child US 19005958

## Publication Classification

**Int. Cl.:** **G06F9/30** (20180101); **G06F9/32** (20180101); **G06F9/38** (20180101)

**U.S. Cl.:**

**CPC** **G06F9/3001** (20130101); **G06F9/30036** (20130101); **G06F9/30038** (20230801); **G06F9/30043** (20130101); **G06F9/3005** (20130101); **G06F9/30076** (20130101); **G06F9/30109** (20130101); **G06F9/30123** (20130101); **G06F9/30145** (20130101); **G06F9/3016** (20130101); **G06F9/323** (20230801); **G06F9/3824** (20130101); **G06F9/383** (20130101);

---

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] The present patent application is a continuation application claiming priority from U.S. patent application Ser. No. 18/449,651 filed Aug. 14, 2024, now U.S. Pat. No. 12,182,568, which is a continuation application claiming priority from U.S. patent application Ser. No. 15/858,916 filed Dec. 29, 2017, now U.S. Pat. No. 11,789,729, each of which is incorporated herein by reference in its entirety.

### FIELD OF INVENTION

[0002] The field of invention relates generally to computer processor architecture, and, more specifically, to systems and methods for computing dot products of nibbles in two tile operands.

### BACKGROUND

[0003] Matrices are increasingly important in many computing tasks such as machine learning and other bulk data processing.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

[0004] The present invention is illustrated by way of example and not limitation in the figures of the accompanying drawings, in which like references indicate similar elements and in which:

[0005] FIG. 1A illustrates an embodiment of configured tiles;

[0006] FIG. 1B illustrates an embodiment of configured tiles;

[0007] FIG. 2 illustrates several examples of matrix storage;

[0008] FIG. 3 illustrates an embodiment of a system utilizing a matrix (tile) operations accelerator;

[0009] FIGS. 4 and 5 show different embodiments of how memory is shared using a matrix operations accelerator;

[0010] FIG. 6 illustrates an embodiment of matrix multiply accumulate operation using tiles (“TMMA”);

[0011] FIG. 7 illustrates an embodiment of a subset of the execution of an iteration of a chained fused multiply accumulate instruction;

[0012] FIG. 8 illustrates an embodiment of a subset of the execution of an iteration of a chained fused multiply accumulate instruction;

[0013] FIG. 9 illustrates an embodiment of a subset of the execution of an iteration of a chained fused multiply accumulate instruction;

[0014] FIG. **10** illustrates an embodiment of a subset of the execution of an iteration of chained fused multiply accumulate instruction;

[0015] FIG. **11** illustrates power-of-two sized SIMD implementations wherein the accumulators use input sizes that are larger than the inputs to the multipliers according to an embodiment;

[0016] FIG. **12** illustrates an embodiment of a system utilizing matrix operations circuitry;

[0017] FIG. **13** illustrates an embodiment of a processor core pipeline supporting matrix operations using tiles;

[0018] FIG. **14** illustrates an embodiment of a processor core pipeline supporting matrix operations using tiles;

[0019] FIG. **15** illustrates an example of a matrix expressed in row major format and column major format;

[0020] FIG. **16** illustrates an example of usage of matrices (tiles);

[0021] FIG. **17** illustrates an embodiment a method of usage of matrices (tiles);

[0022] FIG. **18** illustrates support for configuration of the usage of tiles according to an embodiment;

[0023] FIG. **19** illustrates an embodiment of a description of the matrices (tiles) to be supported;

[0024] FIGS. **20A-20D** illustrate examples of register(s);

[0025] FIG. **21** illustrates an exemplary execution of a TILELOADQPAIR instruction;

[0026] FIG. **22** illustrates an embodiment of a processor executing a TILELOADQPAIR instruction;

[0027] FIG. **23** illustrates a more detailed description of an execution of a TILELOADQPAIR instruction;

[0028] FIG. **24A** is exemplary pseudocode describing an embodiment of a processor executing a TILELOADQPAIR instruction;

[0029] FIG. **24B** is exemplary pseudocode describing an embodiment of helper functions used by a processor executing a TILELOADQPAIR instruction;

[0030] FIGS. **25A-25B** are block diagrams illustrating a generic vector friendly instruction format and instruction templates thereof according to embodiments of the invention;

[0031] FIG. **25A** is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to embodiments of the invention;

[0032] FIG. **25B** is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to embodiments of the invention;

[0033] FIG. **26A** is a block diagram illustrating an exemplary specific vector friendly instruction format according to embodiments of the invention;

[0034] FIG. **26B** is a block diagram illustrating the fields of the specific vector friendly instruction format that make up the full opcode field according to one embodiment of the invention;

[0035] FIG. **26C** is a block diagram illustrating the fields of the specific vector friendly instruction format that make up the register index field according to one embodiment of the invention;

[0036] FIG. **26D** is a block diagram illustrating the fields of the specific vector friendly instruction format that make up the augmentation operation field according to one embodiment of the invention;

[0037] FIG. **27** is a block diagram of a register architecture according to one embodiment of the invention;

[0038] FIG. **28A** is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to embodiments of the invention;

[0039] FIG. **28B** is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to embodiments of the invention;

[0040] FIGS. **29A-B** illustrate a block diagram of a more specific exemplary in-order core

architecture, which core would be one of several logic blocks (including other cores of the same type and/or different types) in a chip;

[0041] FIG. **29A** is a block diagram of a single processor core, along with its connection to the on-die interconnect network and with its local subset of the Level 2 (L2) cache, according to embodiments of the invention;

[0042] FIG. **29B** is an expanded view of part of the processor core in FIG. **29A** according to embodiments of the invention;

[0043] FIG. **30** is a block diagram of a processor that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to embodiments of the invention;

[0044] FIGS. **31-34** are block diagrams of exemplary computer architectures;

[0045] FIG. **31** is a block diagram of a system in accordance with one embodiment of the present invention;

[0046] FIG. **32** is a block diagram of a first more specific exemplary system in accordance with an embodiment of the present invention;

[0047] FIG. **33** is a block diagram of a second more specific exemplary system in accordance with an embodiment of the present invention;

[0048] FIG. **34** is a block diagram of a System-on-a-Chip (SoC) in accordance with an embodiment of the present invention; and

[0049] FIG. **35** is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the invention.

#### DETAILED DESCRIPTION

[0050] In the following description, numerous specific details are set forth. However, it is understood that embodiments of the invention may be practiced without these specific details. In other instances, well-known circuits, structures and techniques have not been shown in detail in order not to obscure the understanding of this description.

[0051] References in the specification to “one embodiment,” “an embodiment,” “an example embodiment,” etc., indicate that the embodiment described may include a particular feature, structure, or characteristic, but every embodiment may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same embodiment. Further, when a particular feature, structure, or characteristic is described in connection with an embodiment, it is submitted that it is within the knowledge of one skilled in the art to affect such feature, structure, or characteristic in connection with other embodiments whether or not explicitly described.

[0052] In many mainstream processors, handling matrices is a difficult and/or instruction intensive task. For example, rows of a matrix could be put into a plurality of packed data (e.g., SIMD or vector) registers and then operated on individually. For example, an add two  $8 \times 2$  matrices may require a load or gather into four packed data registers depending upon data sizes. Then a first add of packed data registers corresponding to a first row from each matrix is performed and a second add of packed data registers corresponding to a second row from each matrix is performed. Then the resulting packed data registers are scattered back to memory. While for small matrices this scenario may be acceptable, it is often not acceptable with larger matrices.

#### I. High-Level Discussion

[0053] Described herein are mechanisms to support matrix operations in computer hardware such as central processing units (CPUs), graphic processing units (GPUs), and accelerators. The matrix operations utilize 2-dimensional (2-D) data structures representing one or more packed regions of memory such as registers. Throughout this description, these 2-D data structures are referred to as tiles. Note that a matrix may be smaller than a tile (use less than all of a tile), or utilize a plurality of tiles (the matrix is larger than the size of any one tile). Throughout the description, matrix (tile)

language is used to indicate operations performed using tiles that impact a matrix; whether or not that matrix is larger than any one tile is not typically relevant.

[0054] Each tile may be acted upon by different operations such as those that are detailed herein and include, but are not limited to: matrix (tile) multiplication, tile add, tile subtract, tile diagonal, tile zero, tile transpose, tile dot product, tile broadcast, tile row broadcast, tile column broadcast, tile multiplication, tile multiplication and accumulation, tile move, etc. Additionally, support for operators such as the use of a scale and/or bias may be used with these operations or in support of non-numeric applications in the future, for instance, OpenCL “local memory,” data compression/decompression, etc.

[0055] Portions of storage (such as memory (non-volatile and volatile), registers, cache, etc.) are arranged into tiles of different horizontal and vertical dimensions. For example, a tile may have horizontal dimension of 4 (e.g., four rows of a matrix) and a vertical dimension of 8 (e.g., 8 columns of the matrix). Typically, the horizontal dimension is related to element sizes (e.g., 2-, 4-, 8-, 16-, 32-, 64-, 128-bit, etc.). Multiple datatypes (single precision floating point, double precision floating point, integer, etc.) may be supported.

#### A. Exemplary Usage of Configured Tiles

[0056] In some embodiments, tile parameters can be configured. For example, a given tile may be configured to provide tile options. Exemplary tile options include, but are not limited to: a number of rows of the tile, a number of columns of the tile, whether the tile is VALID, and whether the tile consists of a PAIR of equal-sized tiles.

[0057] FIG. 1A illustrates an embodiment of configured tiles. As shown, 4 kB of application memory **102** have stored thereon 4 1 kB tiles, tile 0 **104**, tile 1 **106**, tile 2 **108**, and tile 3 **110**. In this example, the 4 tiles do not consist of pairs, and each have elements arranged in rows and columns. Tile t0 **104** and tile t1 **106** have K rows and N columns of 4-byte elements (e.g., single precision data), where K equals 8 and N=32. Tile t2 **108** and tile t3 **110** have K rows and N/2 columns of 8-byte elements (e.g., double precision data). As the double precision operands are twice the width of single precision, this configuration is consistent with a palette, used to provide tile options, supplying at least 4 names with total storage of at least 4 kB. In operation, the tiles can be loaded from and stored to memory using load and store operations. Depending upon the instruction encoding scheme used, the amount of available application memory, as well as the size, number, and configuration of available tiles varies.

[0058] FIG. 1B illustrates an embodiment of configured tiles. As shown, 4 kB of application memory **122** have stored thereon 2 pairs of 1 kB-tiles, the first pair being tile t4L **124** and tile t4R **126**, and the second pair being tile t5L **128** and tile t5R **130**. As shown the pairs of tiles are divided into a left tile and a right tile. In other embodiments, the pair of tiles are divided into an even tile and an odd tile. In this example, the 4 tiles each have elements arranged in rows and columns. Tile t4L **124** and tile t4R **126** have K rows and N columns of 4-byte elements (e.g., single precision data), where K equals 8 and N equals 32. Tile t5L **128** and tile t5R **130** have K rows and N/2 columns of 8-byte elements (e.g., double precision data). As the double precision operands are twice the width of single precision, this configuration is consistent with a palette, used to provide tile options, supplying at least 2 names with total storage of at least 4 kB. The four tiles of FIG. 1A use 4 names, each naming a 1 kB tile, whereas the 2 pairs of tiles in FIG. 1B can use 2 names to specify the paired tiles. In some embodiments, tile instructions accept a name of a paired tile as an operand. In operation, the tiles can be loaded from and stored to memory using load and store operations. Depending upon the instruction encoding scheme used, the amount of available application memory, as well as the size, number, and configuration of available tiles varies.

[0059] In some embodiments, tile parameters are definable. For example, a “palette” is used to provide tile options. Exemplary options include, but are not limited to: the number of tile names, the number of bytes in a row of storage, the number of rows and columns in a tile, etc. For example, a maximum “height” (number of rows) of a tile may be defined as:

Tile Max Rows=Architected Storage/(The Number of Palette Names\*The Number of Bytes per row).

[0060] As such, an application can be written such that a fixed usage of names will be able to take advantage of different storage sizes across implementations.

[0061] Configuration of tiles is done using a tile configuration (“TILECONFIG”) instruction, where a particular tile usage is defined in a selected palette. This declaration includes the number of tile names to be used, the requested number of rows and columns per name (tile), and, in some embodiments, the requested datatype of each tile. In some embodiments, consistency checks are performed during the execution of a TILECONFIG instruction to determine that it matches the restrictions of the palette entry.

#### B. Exemplary Tile Storage Types

[0062] FIG. 2 illustrates several examples of matrix storage. In (A), a tile is stored in memory. As shown, each “row” consists of four packed data elements. To get to the next “row,” a stride value is used. Note that rows may be consecutively stored in memory. Strided memory accesses allows for access of one row to then next when the tile storage does not map the underlying memory array row width.

[0063] Tile loads from memory and stores to memory are typically strided accesses from the application memory to packed rows of data. Exemplary TILELOAD and TILESTORE instructions, or other instruction references to application memory as a TILE operand in load-op instructions, are, in some embodiments, restartable to handle (up to) 2\*rows of page faults, unmasked floating point exceptions, and/or interrupts per instruction.

[0064] In (B), a matrix is stored in a tile comprised of a plurality of registers such as packed data registers (single instruction, multiple data (SIMD) or vector registers). In this example, the tile is overlaid on three physical registers. Typically, consecutive registers are used, however, this need not be the case.

[0065] In (C), a matrix is stored in a tile in non-register storage accessible to a fused multiple accumulate (FMA) circuit used in tile operations. This storage may be inside of an FMA, or adjacent to it. Additionally, in some embodiments, discussed below, the storage may be for a data element and not an entire row or tile.

[0066] The supported parameters for the TMMA architecture are reported via CPUID. In some embodiments, the list of information includes a maximum height and a maximum SIMD dimension. Configuring the TMMA architecture requires specifying the dimensions for each tile, the element size for each tile and the palette identifier. This configuration is done by executing the TILECONFIG instruction.

[0067] Successful execution of a TILECONFIG instruction enables subsequent TILE operators. A TILERELASEALL instruction clears the tile configuration and disables the TILE operations (until the next TILECONFIG instructions executes). In some embodiments, XSAVE, XSTORE, etc. are used in context switching using tiles. In some embodiments, 2 XCR0 bits are used in XSAVE, one for TILECONFIF metadata and one bit corresponding to actual tile payload data.

[0068] TILECONFIG not only configures the tile usage, but also sets a state variable indicating that the program is in a region of code with tiles configured. An implementation may enumerate restrictions on other instructions that can be used with a tile region such as no usage of an existing register set, etc.

[0069] Exiting a tile region is typically done with the TILERELASEALL instruction. It takes no parameters and swiftly invalidates all tiles (indicating that the data no longer needs any saving or restoring) and clears the internal state corresponding to being in a tile region.

[0070] In some embodiments, tile operations will zero any rows and any columns beyond the dimensions specified by the tile configuration. For example, tile operations will zero the data beyond the configured number of columns (factoring in the size of the elements) as each row is

written. For example, with 64 byte rows and a tile configured with 10 rows and 12 columns, an operation writing FP32 elements would write each of the first 10 rows with 12\*4 bytes with output/result data and zero the remaining 4\*4 bytes in each row. Tile operations also fully zero any rows after the first 10 configured rows. When using 1K tile with 64 byte rows, there would be 16 rows, so in this example, the last 6 rows would also be zeroed.

[0071] In some embodiments, a context restore (e.g., XRSTOR), when loading data, enforces that the data beyond the configured rows for a tile will be maintained as zero. If there is no valid configuration, all rows are zeroed. XRSTOR of tile data can load garbage in the columns beyond those configured. It should not be possible for XRSTOR to clear beyond the number of columns configured because there is not an element width associated with the tile configuration.

[0072] Context save (e.g., XSAVE) exposes the entire TILE storage area when writing it to memory. If XRSTOR loaded garbage data in to the rightmost part of a tile, that data will be saved by XSAVE. XSAVE will write zeros for rows beyond the number specified for each tile.

[0073] In some embodiments, tile instructions are restartable. The operations that access memory allow restart after page faults. The computational instructions that deal with floating point operations also allow for unmasked floating point exceptions, with the masking of the exceptions controlled by a control and/or status register.

[0074] To support restarting instructions after these events, the instructions store information in the start registers detailed below.

## II. Matrix (Tile) Operation Systems

### A. Exemplary Hardware Support

[0075] FIG. 3 illustrates an embodiment of a system utilizing a matrix (tile) operations accelerator. In this illustration, a host processor/processing system 301 communicates commands 311 (e.g., matrix manipulation operations such as arithmetic or matrix manipulation operations, or load and store operations) to a matrix operations accelerator 307. However, this is shown this way for discussion purposes only. As detailed later, this accelerator 307 may be a part of a processing core. Typically, commands 311 that are tile manipulation operator instructions will refer to tiles as register-register (“reg-reg”) or register-memory (“reg-mem”) format. Other commands such as TILESTORE, TILELOAD, TILECONFIG, etc., do not perform data operations on a tile. Commands may be decoded instructions (e.g., micro-ops) or macro-instructions for the accelerator 307 to handle.

[0076] In this example, a coherent memory interface 303 is coupled to the host processor/processing system 301 and matrix operations accelerator 307 such that they can share memory. FIGS. 4 and 5 show different embodiments of how memory is shared using a matrix operations accelerator. As shown in FIG. 4, the host processor 401 and matrix operations accelerator circuitry 405 share the same memory 403. FIG. 5 illustrates an embodiment where the host processor 501 and matrix operations accelerator 505 do not share memory, but can access each other's memory. For example, processor 501 can access tile memory 507 and utilize its host memory 503 as normal. Similarly, the matrix operations accelerator 505 can access host memory 503, but more typically uses its own memory 507. Note these memories may be of different types.

[0077] In some embodiments, the matrix operations accelerator 307 includes a plurality of FMAs 309 coupled to data buffers 305 (in some implementations, one or more of these buffers 305 are stored in the FMAs of the grid as shown). The data buffers 305 buffer tiles loaded from memory and/or tiles to be stored to memory (e.g., using a tileload or tilestore instruction). Data buffers may be, for example, a plurality of registers. Typically, these FMAs are arranged as a grid of chained FMAs 309 which are able to read and write tiles. In this example, the matrix operations accelerator 307 is to perform a matrix multiply operation using tiles T0, T1, and T2. At least one of tiles is housed in the FMA grid 309. In some embodiments, all tiles in an operation are stored in the FMA grid 309. In other embodiments, only a subset is stored in the FMA grid 309. As shown, T1 is housed and T0 and T2 are not. Note that A, B, and C refer to the matrices of these tiles which may

or may not take up the entire space of the tile.

[0078] FIG. 6 illustrates an embodiment of matrix multiply accumulate operation using tiles (“TMMA”).

[0079] The number of rows in the matrix (TILE A **601**) matches the number of serial (chained) FMAs comprising the computation's latency. An implementation is free to recirculate on a grid of smaller height, but the computation remains the same.

[0080] The source/destination vector comes from a tile of N rows (TILE C **605**) and the grid of FMAs **611** performs N vector-matrix operations resulting in a complete instruction performing a matrix multiplication of tiles. Tile B **603** is the other vector source and supplies “broadcast” terms to the FMAs in each stage.

[0081] In operation, in some embodiments, the elements of matrix B (stored in a tile B **603**) are spread across the rectangular grid of FMAs. Matrix B (stored in tile A **601**) has its elements of a row transposed to match up with the columnar dimension of the rectangular grid of FMAs. At each FMA in the grid, an element of A and B are multiplied and added to the incoming summand (from above in the Figure) and the outgoing sum is passed to the next row of FMAs (or the final output).

[0082] The latency of a single step is proportional to K (row height of matrix B) and dependent TMMA typically have enough source-destination rows (either in a single tile or across tile) to hide that latency. An implementation may also split the SIMD (packed data element) dimension M (row height of matrix A) across time steps, but this simply changes the constant that K is multiplied by. When a program specifies a smaller K than the maximum enumerated by the TMACC, an implementation is free to implement this with “masking” or “early outs.”

[0083] The latency of an entire TMMA is proportional to  $N \cdot K$ . The repeat rate is proportional to N. The number of MACs per TMMA instruction is  $N \cdot K \cdot M$ .

[0084] FIG. 7 illustrates an embodiment of a subset of the execution of an iteration of a chained fused multiply accumulate instruction. In particular, this illustrates execution circuitry of an iteration of one packed data element position of the destination. In this embodiment, the chained fused multiply accumulate is operating on signed sources wherein the accumulator is  $2 \times$  the input data size.

[0085] A first signed source (source 1 **701**) and a second signed source (source 2 **703**) each have four packed data elements. Each of these packed data elements stores signed data such as floating point data. A third signed source (source 3 **709**) has two packed data elements, each of which stores signed data. The sizes of the first and second signed sources **701** and **703** are half that of the third signed source (initial value or previous result) **709**. For example, the first and second signed sources **701** and **703** could have 32-bit packed data elements (e.g., single precision floating point) while the third signed source **709** could have 64-bit packed data elements (e.g., double precision floating point).

[0086] In this illustration, only the two most significant packed data element positions of the first and second signed sources **701** and **703** and the most significant packed data element position of the third signed source **709** are shown. Of course, the other packed data element positions would also be processed.

[0087] As illustrated, packed data elements are processed in pairs. For example, the data of the most significant packed data element positions of the first and second signed sources **701** and **703** are multiplied using a multiplier circuit **705**, and the data from second most significant packed data element positions of the first and second signed sources **701** and **703** are multiplied using a multiplier circuit **707**. In some embodiments, these multiplier circuits **705** and **707** are reused for other packed data elements positions. In other embodiments, additional multiplier circuits are used so that the packed data elements are processed in parallel. In some contexts, parallel execution is done using lanes that are the size of the signed third source **709**. The results of each of the multiplications are added using addition circuitry **711**.

[0088] The result of the addition of the results of the multiplications is added to the data from most



significant packed data element position of the signed source 3 **709** (using a different adder **713** or the same adder **711**).

[0089] Finally, the result of the second addition is either stored into the signed destination **715** in a packed data element position that corresponds to the packed data element position used from the signed third source **709**, or passed on to the next iteration, if there is one. In some embodiments, a writemask is applied to this storage such that if a corresponding writemask (bit) is set, the storage happens, and, if not set, the storage does not happen.

[0090] FIG. **8** illustrates an embodiment of a subset of the execution of an iteration of a chained fused multiply accumulate instruction. In particular, this illustrates execution circuitry of an iteration of one packed data element position of the destination. In this embodiment, the chained fused multiply accumulate is operating on signed sources wherein the accumulator is  $2\times$  the input data size.

[0091] A first signed source (source 1 **801**) and a second signed source (source 2 **803**) each have four packed data elements. Each of these packed data elements stores signed data such as integer data. A third signed source (source 3 **809**) has two packed data elements, each of which stores signed data. The sizes of the first and second signed sources **801** and **803** are half that of the third signed source **809**. For example, the first and second signed sources **801** and **803** could have 32-bit packed data elements (e.g., single precision floating point) the third signed source **809** could have 64-bit packed data elements (e.g., double precision floating point).

[0092] In this illustration, only the two most significant packed data element positions of the first and second signed sources **801** and **803** and the most significant packed data element position of the third signed source **809** are shown. Of course, the other packed data element positions would also be processed.

[0093] As illustrated, packed data elements are processed in pairs. For example, the data of the most significant packed data element positions of the first and second signed sources **801** and **803** are multiplied using a multiplier circuit **805**, and the data from second most significant packed data element positions of the first and second signed sources **801** and **803** are multiplied using a multiplier circuit **807**. In some embodiments, these multiplier circuits **805** and **807** are reused for other packed data elements positions. In other embodiments, additional multiplier circuits are used so that the packed data elements are processed in parallel. In some contexts, parallel execution is done using lanes that are the size of the signed third source (initial value or previous iteration result) **809**. The results of each of the multiplications are added to the signed third source **809** using addition/saturation circuitry **813**.

[0094] Addition/saturation (accumulator) circuitry **813** preserves a sign of an operand when the addition results in a value that is too big. In particular, saturation evaluation occurs on the infinite precision result between the multi-way-add and the write to the destination or next iteration. When the accumulator **813** is floating point and the input terms are integer, the sum of products and the floating point accumulator input value are turned into infinite precision values (fixed point numbers of hundreds of bits), the addition of the multiplication results and the third input is performed, and a single rounding to the actual accumulator type is performed.

[0095] Unsigned saturation means the output values are limited to a maximum unsigned number for that element width (all 1s). Signed saturation means a value is limited to be in the range between a minimum negative number and a max positive number for that element width (for bytes for example, the range is from  $-128 (= -2^{\lceil \log_2(256) \rceil - 1})$  to  $127 (= 2^{\lceil \log_2(256) \rceil - 1} - 1)$ ).

[0096] The result of the addition and saturation check is stored into the signed result **815** in a packed data element position that corresponds to the packed data element position used from the signed third source **809**, or passed on to the next iteration if there is one. In some embodiments, a writemask is applied to this storage such that if a corresponding writemask (bit) is set, the storage happens, and, if not set, the storage does not happen.

[0097] FIG. 9 illustrates an embodiment of a subset of the execution of an iteration of a chained fused multiply accumulate instruction. In particular, this illustrates execution circuitry of an iteration of one packed data element position of the destination. In this embodiment, the chained fused multiply accumulate is operating on a signed source and an unsigned source wherein the accumulator is  $4\times$  the input data size.

[0098] A first signed source (source 1 **901**) and a second unsigned source (source 2 **903**) each have four packed data elements. Each of these packed data elements has data such as floating point or integer data. A third signed source (initial value or result **915**) has a packed data element of which stores signed data. The sizes of the first and second sources **901** and **903** are a quarter of the third signed source **915**. For example, the first and second sources **901** and **903** could have 16-bit packed data elements (e.g., word) and the third signed source **915** could have 64-bit packed data elements (e.g., double precision floating point or 64-bit integer).

[0099] In this illustration, the four most significant packed data element positions of the first and second sources **901** and **903** and the most significant packed data element position of the third signed source **915** are shown. Of course, other packed data element positions would also be processed if there are any

[0100] As illustrated, packed data elements are processed in quadruplets. For example, the data of the most significant packed data element positions of the first and second sources **901** and **903** are multiplied using a multiplier circuit **905**, data from second most significant packed data element positions of the first and second sources **901** and **903** are multiplied using a multiplier circuit **907**, data from third most significant packed data element positions of the first and second sources **901** and **903** are multiplied using a multiplier circuit **909**, and data from the least significant packed data element positions of the first and second sources **901** and **903** are multiplied using a multiplier circuit **911**. In some embodiments, the signed packed data elements of the first source **901** are sign extended and the unsigned packed data elements of the second source **903** are zero extended prior to the multiplications.

[0101] In some embodiments, these multiplier circuits **905-911** are reused for other packed data elements positions. In other embodiments, additional multiplier circuits are used so that the packed data elements are processed in parallel. In some contexts, parallel execution is done using lanes that are the size of the signed third source **915**. The results of each of the multiplications are added using addition circuitry **911**.

[0102] The result of the addition of the results of the multiplications is added to the data from most significant packed data element position of the signed source 3 **915** (using a different adder **913** or the same adder **911**).

[0103] Finally, the result **919** of the second addition is either stored into the signed destination in a packed data element position that corresponds to the packed data element position used from the signed third source **915**, or passed to the next iteration. In some embodiments, a writemask is applied to this storage such that if a corresponding writemask (bit) is set, the storage happens, and, if not set, the storage does not happen.

[0104] FIG. 10 illustrates an embodiment of a subset of the execution of an iteration of chained fused multiply accumulate instruction. In particular, this illustrates execution circuitry of an iteration of one packed data element position of the destination. In this embodiment, the chained fused multiply accumulate is operating on a signed source and an unsigned source wherein the accumulator is  $4\times$  the input data size.

[0105] A first signed source (source 1 **1001**) and a second unsigned source (source 2 **1003**) each have four packed data elements. Each of these packed data elements stores data such as floating point or integer data. A third signed source (initial or previous result **1015**) has a packed data element of which stores signed data. The sizes of the first and second sources **1001** and **1003** are a quarter of the third signed source 1 **015**. For example, the first and second sources **1001** and **1003** could have 16-bit packed data elements (e.g., word) and the third signed source 1 **015** could have

64-bit packed data elements (e.g., double precision floating point or 64-bit integer).

[0106] In this illustration, the four most significant packed data element positions of the first and second sources **1001** and **1003** and the most significant packed data element position of the third signed source 1 **015** are shown. Of course, other packed data element positions would also be processed if there are any.

[0107] As illustrated, packed data elements are processed in quadruplets. For example, the data of the most significant packed data element positions of the first and second sources **1001** and **1003** are multiplied using a multiplier circuit **1005**, data from second most significant packed data element positions of the first and second sources **1001** and **1003** are multiplied using a multiplier circuit **1007**, data from third most significant packed data element positions of the first and second sources **1001** and **1003** are multiplied using a multiplier circuit **1009**, and data from the least significant packed data element positions of the first and second sources **1001** and **1003** are multiplied using a multiplier circuit **1011**. In some embodiments, the signed packed data elements of the first source 1 **001** are sign extended and the unsigned packed data elements of the second source 1 **003** are zero extended prior to the multiplications.

[0108] In some embodiments, these multiplier circuits **1005-1011** are reused for other packed data elements positions. In other embodiments, additional multiplier circuits are used so that the packed data elements are processed in parallel. In some contexts, parallel execution is done using lanes that are the size of the signed third source 1 **015**. The result of the addition of the results of the multiplications is added to the data from most significant packed data element position of the signed source 3 **1015** using addition/saturation circuitry **1013**.

[0109] Addition/saturation (accumulator) circuitry **1013** preserves a sign of an operand when the addition results in a value that is too big or too small for signed saturation. In particular, saturation evaluation occurs on the infinite precision result between the multi-way-add and the write to the destination. When the accumulator **1013** is floating point and the input terms are integer, the sum of products and the floating point accumulator input value are turned into infinite precision values (fixed point numbers of hundreds of bits), the addition of the multiplication results and the third input is performed, and a single rounding to the actual accumulator type is performed.

[0110] The result **1019** of the addition and saturation check is stored into the signed destination in a packed data element position that corresponds to the packed data element position used from the signed third source 1 **015**, or passed to the next iteration. In some embodiments, a writemask is applied to this storage such that if a corresponding writemask (bit) is set, the storage happens, and, if not set, the storage does not happen.

[0111] FIG. **11** illustrates power-of-two sized SIMD implementations wherein the accumulators use input sizes that are larger than the inputs to the multipliers according to an embodiment. Note the source (to the multipliers) and accumulator values may be signed or unsigned values. For an accumulator having  $2\times$  input sizes (in other words, the accumulator input value is twice the size of the packed data element sizes of the sources), table **1101** illustrates different configurations. For byte sized sources, the accumulator uses word or half-precision floating-point (HPFP) values that are 16-bit in size. For word sized sources, the accumulator uses 32-bit integer or single-precision floating-point (SPFP) values that are 32-bit in size. For SPFP or 32-bit integer sized sources, the accumulator uses 64-bit integer or double-precision floating-point (DPFP) values that are 64-bit in size.

[0112] For an accumulator having  $4\times$  input sizes (in other words, the accumulator input value is four times the size of the packed data element sizes of the sources), table **1103** illustrates different configurations. For byte sized sources, the accumulator uses 32-bit integer or single-precision floating-point (SPFP) values that are 32-bit in size. For word sized sources, the accumulator uses 64-bit integer or double-precision floating-point (DPFP) values that are 64-bit in size in some embodiments.

[0113] For an accumulator having  $8\times$  input sizes (in other words, the accumulator input value is

eight times the size of the packed data element sizes of the sources), table **1105** illustrates a configuration. For byte sized sources, the accumulator uses 64-bit integer.

[0114] As hinted at earlier, matrix operations circuitry may be included in a core, or as an external accelerator. FIG. **12** illustrates an embodiment of a system utilizing matrix operations circuitry. In this illustration, a plurality of entities are coupled with a ring interconnect **1245**.

[0115] A plurality of cores **1201**, **1203**, **1205**, and **1207** provide non-tile based instruction support. In some embodiments, matrix operations circuitry **1251** is provided in a core **1203**, and in other embodiments matrix operations circuitry **1211** and **1213** are accessible on the ring interconnect **1245**.

[0116] Additionally, one or more memory controllers **1223-1225** are provided to communicate with memory **1233** and **1231** on behalf of the cores and/or matrix operations circuitry.

[0117] FIG. **13** illustrates an embodiment of a processor core pipeline supporting matrix operations using tiles. Branch prediction and decode circuitry **1303** performs branch predicting of instructions, decoding of instructions, and/or both from instructions stored in instruction storage **1301**. For example, instructions detailed herein may be stored in instruction storage. In some implementations, separate circuitry is used for branch prediction and in some embodiments, at least some instructions are decoded into one or more micro-operations, micro-code entry points, microinstructions, other instructions, or other control signals using microcode **1305**. The branch prediction and decode circuitry **1303** may be implemented using various different mechanisms.

Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc. [0118] The branch prediction and decode circuitry **1303** is coupled to a rename/allocator circuitry **1307** which is coupled, in some embodiments, to scheduler circuitry **1309**. In some embodiments, these circuits provide register renaming, register allocation, and/or scheduling functionality by performing one or more of: 1) renaming logical operand values to physical operand values (e.g., a register alias table in some embodiments), 2) allocating status bits and flags to the decoded instruction, and 3) scheduling the decoded instruction for execution on execution circuitry out of an instruction pool (e.g., using a reservation station in some embodiments).

[0119] The scheduler circuitry **1309** represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) scheduler circuitry **1309** is coupled to, or includes, physical register file(s) **1315**. Each of the physical register file(s) **1315** represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point, status (e.g., an instruction pointer that is the address of the next instruction to be executed), tiles, etc. In one embodiment, the physical register file(s) **1315** comprises vector registers circuitry, write mask registers circuitry, and scalar registers circuitry. These register circuits may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) **1315** is overlapped by a retirement circuit **1317** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement circuit **1317** and the physical register file(s) **1315** are coupled to the execution circuit(s) **1311**.

[0120] While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used in an in-order architecture. While the illustrated embodiment of the processor may also include separate instruction and data cache units and a shared L2 cache unit, alternative embodiments may have a single internal cache for both instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some embodiments, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache

may be external to the core and/or the processor.

[0121] The execution circuitry **1311** is a set of one or more execution circuits **1321**, **1323**, and **1327** and a set of one or more memory access circuits **1325**. The execution circuits **1321**, **1323**, and **1327** perform various operations (e.g., shifts, addition, subtraction, multiplication) and on various types of data (e.g., scalar floating point, packed integer, packed floating point, vector integer, vector floating point). While some embodiments may include a number of execution units dedicated to specific functions or sets of functions, other embodiments may include only one execution unit or multiple execution units that all perform all functions. The scalar circuitry **1321** performs scalar operations, the vector/SIMD circuitry **1323** performs vector/SIMD operations, and matrix operations circuitry **1327** performs matrix (tile) operations detailed herein.

[0122] By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement a pipeline as follows: 1) an instruction fetch circuit performs fetch and length decoding stages; 2) the branch and decode circuitry **1303** performs a decode stage; 3) the rename/allocator circuitry **1307** performs an allocation stage and renaming stage; 4) the scheduler circuitry **1309** performs a schedule stage; 5) physical register file(s) (coupled to, or included in, the scheduler circuitry **1309** and rename/allocate circuitry **1307** and a memory unit perform a register read/memory read stage; the execution circuitry **1311** performs an execute stage; 6) a memory unit and the physical register file(s) unit(s) perform a write back/memory write stage; 7) various units may be involved in the exception handling stage; and 8) a retirement unit and the physical register file(s) unit(s) perform a commit stage.

[0123] The core may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, CA; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, CA), including the instruction(s) described herein. In one embodiment, the core **1390** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

[0124] It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyperthreading technology).

[0125] FIG. **14** illustrates an embodiment of a processor core pipeline supporting matrix operations using tiles. Branch prediction and decode circuitry **1403** performs branch predicting of instructions, decoding of instructions, and/or both from instructions stored in instruction storage **1401**. For example, instructions detailed herein may be stored in instruction storage. In some implementations, separate circuitry is used for branch prediction and in some embodiments, at least some instructions are decoded into one or more micro-operations, micro-code entry points, microinstructions, other instructions, or other control signals using microcode **1405**. The branch prediction and decode circuitry **1403** may be implemented using various different mechanisms. Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc.

[0126] The branch prediction and decode circuitry **1403** is coupled to a rename/allocator circuitry **1407** which is coupled, in some embodiments, to scheduler circuitry **1409**. In some embodiments, these circuits provide register renaming, register allocation, and/or scheduling functionality by performing one or more of: 1) renaming logical operand values to physical operand values (e.g., a register alias table in some embodiments), 2) allocating status bits and flags to the decoded instruction, and 3) scheduling the decoded instruction for execution on execution circuitry out of an instruction pool (e.g., using a reservation station in some embodiments).

[0127] The scheduler circuitry **1409** represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) scheduler circuitry **1409** is coupled to, or includes, physical register file(s) **1415**. Each of the physical register file(s) **1415** represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point, status (e.g., an instruction pointer that is the address of the next instruction to be executed), tiles, etc. In one embodiment, the physical register file(s) **1415** comprises vector registers circuitry, write mask registers circuitry, and scalar registers circuitry. These register circuits may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) **1415** is overlapped by a retirement circuit **1417** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement circuit **1417** and the physical register file(s) **1415** are coupled to the execution circuit(s) **1411**.

[0128] While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used in an in-order architecture. While the illustrated embodiment of the processor may also include separate instruction and data cache units and a shared L2 cache unit, alternative embodiments may have a single internal cache for both instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some embodiments, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache may be external to the core and/or the processor.

[0129] The execution circuitry **1411** a set of one or more execution circuits **1427** and a set of one or more memory access circuits **1425**. The execution circuits **1427** perform matrix (tile) operations detailed herein.

[0130] By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement a pipeline as follows: 1) an instruction fetch circuit performs fetch and length decoding stages; 2) the branch and decode circuitry **1403** performs a decode stage; 3) the rename/allocator circuitry **1407** performs an allocation stage and renaming stage; 4) the scheduler circuitry **1409** performs a schedule stage; 5) physical register file(s) (coupled to, or included in, the scheduler circuitry **1407** and rename/allocate circuitry **1407** and a memory unit perform a register read/memory read stage; the execution circuitry **1411** performs an execute stage; 6) a memory unit and the physical register file(s) unit(s) perform a write back/memory write stage; 7) various units may be involved in the exception handling stage; and 8) a retirement unit and the physical register file(s) unit(s) perform a commit stage.

[0131] The core may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, CA; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, CA), including the instruction(s) described herein. In one embodiment, the core **1490** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

[0132] It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyperthreading technology).

B. Layout

[0133] Throughout this description, data is expressed using row major data layout. Column major users should translate the terms according to their orientation. FIG. 15 illustrates an example of a matrix expressed in row major format and column major format. As shown, matrix A is a  $2 \times 3$  matrix. When this matrix is stored in row major format, the data elements of a row are consecutive. When this matrix is stored in column major format, the data elements of a column are consecutive. It is a well-known property of matrices that  $A^{sup.T} \cdot B^{sup.T} = (BA)^{sup.T}$ , where superscript T means transpose. Reading column major data as row major data results in the matrix looking like the transpose matrix.

[0134] In some embodiments, row-major semantics are utilized in hardware, and column major data is to swap the operand order with the result being transposes of matrix, but for subsequent column-major reads from memory it is the correct, non-transposed matrix.

[0135] For example, if there are two column-major matrices to multiply:

TABLE-US-00001  $a \ b \ g \ i \ k \ ag + bh \ ai + bj \ ak + bl \ c \ d * h \ j \ l = cg + dh \ ci + dj \ ck + dl \ e \ f \ eg + fh \ ei + fj \ ek + fl \ (3 \times 2) \ (2 \times 3) \ (3 \times 3)$

[0136] The input matrices would be stored in linear memory (column-major) as: [0137]  $a \ c \ e \ b \ d \ f$  [0138] and [0139]  $g \ h \ i \ j \ k \ l$ .

[0140] Reading those matrices as row-major with dimensions  $2 \times 3$  and  $3 \times 2$ , they would appear as: TABLE-US-00002  $a \ c \ e$  and  $g \ h \ b \ d \ f \ i \ j \ k \ l$

[0141] Swapping the order and matrix multiplying:

TABLE-US-00003  $g \ h \ a \ c \ e \ ag + bh \ cg + dh \ eg + fh \ i \ j * b \ d \ f = ai + bj \ ci + dj \ ei + fj \ k \ l \ ak + bl \ ck + dl \ ek + fl$  [0142] the transpose matrix is out and can then be stored in in row-major order: [0143]  $ag + bh \ cg + dh \ eg + fh \ ai + bj \ ci + dj \ ei + fj \ ak + bl \ ck + dl \ ek + fl$  [0144] and used in subsequent column major computations, it is the correct un-transposed matrix:

TABLE-US-00004  $ag + bh \ ai + bj \ ak + bl \ cg + dh \ ci + dj \ ck + dl \ eg + fh \ ei + fj \ ek + fl$

### III. Exemplary Usage

[0145] FIG. 16 illustrates an example of usage of matrices (tiles). In this example, matrix C **1601** includes two tiles, matrix A **1603** includes one tile, and matrix B **1605** includes two tiles. This figure shows an example of the inner loop of an algorithm to compute a matrix multiplication. In this example, two result tiles, tmm0 and tmm1, from matrix C **1601** are used to accumulate the intermediate results. One tile from the A matrix **1603** (tmm2) is re-used twice as it multiplied by two tiles from the B matrix **1605**. Pointers to load a new A tile and two new B tiles from the directions indicated by the arrows. An outer loop, not shown, adjusts the pointers for the C tiles. [0146] The exemplary code as shown includes the usage of a tile configuration instruction and is executed to configure tile usage, load tiles, a loop to process the tiles, store tiles to memory, and release tile usage.

[0147] FIG. 17 illustrates an embodiment of usage of matrices (tiles). At **1701**, tile usage is configured. For example, a TILECONFIG instruction is executed to configure tile usage including setting a number of rows and columns per tile. Typically, at least one matrix (tile) is loaded from memory at **1703**. At least one matrix (tile) operation is performed at **1705** using the matrices (tiles). At **1707**, at least one matrix (tile) is stored out to memory and a context switch can occur at **1709**.

### IV. Exemplary Configuration

#### A. Tile Configuration Hardware Support

[0148] As discussed above, tile usage typically needs to be configured prior to use. For example, full usage of all rows and columns may not be needed. Not only does not configuring these rows and columns save power in some embodiments, but the configuration may be used to determine if an operation will generate an error. For example, a matrix multiplication of the form  $(N \times M) * (L \times N)$  will typically not work if M and L are not the same.

[0149] Prior to using matrices using tiles, in some embodiments, tile support is to be configured. For example, how many rows and columns per tile, tiles that are to be used, etc. are configured. A TILECONFIG instruction is an improvement to a computer itself as it provides for support to

configure the computer to use a matrix accelerator (either as a part of a processor core, or as an external device). In particular, an execution of the TILECONFIG instruction causes a configuration to be retrieved from memory and applied to matrix (tile) settings within a matrix accelerator.

#### i. Tile Usage Configuration

[0150] FIG. **18** illustrates support for configuration of the usage of tiles according to an embodiment. A memory **1801** contains the description of the matrices (tiles) to be supported **1803**.

[0151] Execution circuitry **1811** of a processor/core **1805** stores aspects of a tile description **1803** into tile configurations **1817**. The tile configurations **1817** detail what tiles for a palette are configured (the number of rows and columns in each tile) and a marking that matrix support is in use. In particular, instruction execution resources **1811** are configured to use tiles as specified by the tile configuration **1817**. The instruction execution resources may also include a machine specific register or configuration register to indicate tile usage. Additional values such as in-use and start values are also set. The tile configurations **1817** utilize one or more registers **1819** to store tile usage and configuration information.

[0152] FIG. **19** illustrates an embodiment of a description of the matrices (tiles) to be supported. This is the description that is to be stored upon an execution of a STTILECFG instruction. In this example, each field is a byte. In byte[0], a palette ID **1901** is stored. The palette ID is used to index a palette table **1813** which stores, per palette ID, a number of bytes in a tile, and bytes per row of the tiles that are associated with this ID as defined by the configuration.

[0153] Byte 1 stores a value to be stored in a “startRow” register **1903** and byte 2 stores a value to be stored in a “startP” register **1905**. To support restarting instructions after these events, the instructions store information these registers. To support restarting instructions after break events such as those detailed above, the instructions store information in these registers. The startRow value indicates the row that should be used for restart. The startP value indicates the position within the row for store operations when pairs are used and, in some embodiments, indicates the lower half of the row (in the lower tile of a pair) or higher half of the row (in the higher tile of a pair). Generally, the position in the row (the column) is not needed.

[0154] With the exception of TILECONFIG and STTILECFG, successfully executing matrix (tile) instructions will set both startRow and startP to zero.

[0155] Any time an interrupted matrix (tile) instruction is not restarted, it is the responsibility of software to zero the startRow and startP values. For example, unmasked floating point exception handlers might decide to finish the operation in software and change the program counter value to another instruction, usually the next instruction. In this case the software exception handler must zero the startRow and startP values in the exception presented to it by the operating system before resuming the program. The operating system will subsequently reload those values using a restore instruction.

[0156] Byte 3 stores an indication of pairs (**1b** per tile) of tiles **1907**.

[0157] Bytes 16-17 store the number of rows **1913** and columns **1915** for tile 0, bytes 18-19 store the number of rows and columns for tile 1, etc. In other words, each 2 byte group specifies a number of rows and columns for a tile. If a group of 2 bytes is not used to specify tile parameters, they should have the value zero. Specifying tile parameters for more tiles than the implementation limit or the palette limit results in a fault. Unconfigured tiles are set to an initial state with 0 rows, 0 columns.

[0158] Finally, the configuration in memory typically ends with an ending delineation such as all zeros for several consecutive bytes.

#### ii. Exemplary Tile and Tile Configuration Storage

[0159] FIGS. **20A-20D** illustrate examples of register(s) **1819**. FIG. **20A** illustrates a plurality of registers **1819**. As shown each tile (TMM0 **2001** . . . TMMN **2003**) has a separate register with each register storing a row and column size for that particular tile. StartP and StartRow are stored in separate registers **2011** and **2013**. One or more status registers **2015** are set (e.g.,



TILES\_CONFIGURED=1) to indicate tiles are configured for use.

[0160] FIG. 20B illustrates a plurality of registers **1819**. As shown each tile has separate registers for its rows and columns. For example, TMM0 rows configuration **2021**, TMM0 columns configuration **2023**, StartP and StartRow are stored in separate registers **2011** and **2013**. One or more status registers **2015** are set (e.g., TILES\_CONFIGURED=1) to indicate tiles are configured for use.

[0161] FIG. 20C illustrates a single register **1819**. As shown, this register stores tile configurations (rows and columns per tile) **2031**, StartP **2011**, and StartRow **2013** are stored in single register as packed data registers. One or more status registers **2015** are set (e.g., TILES\_CONFIGURED=1) to indicate tiles are configured for use.

[0162] FIG. 20D illustrates a plurality of registers **1819**. As shown, a single register stores tile configurations (rows and columns per tile) **2031**. StartP and StartRow are stored in separate registers **2011** and **2013**. One or more status registers **2015** are set (e.g., TILES\_CONFIGURED=1) to indicate tiles are configured for use.

[0163] Other combinations are contemplated such as combining the start registers into a single register where they are shown separately, etc.

## B. TDP4BIT

[0164] A common matrix operation is a dot product, for example between elements of two tiles tile stored in a set of registers or in memory. Detailed herein are embodiments of a TDP4BIT instruction and its execution. A TDP4BIT instruction is an improvement to a computer itself as it provides for support to perform a dot product between elements of two matrices within a matrix (tile) with a single instruction. In particular, the execution of the TDP4BIT instruction performs a dot product between the elements of a first source matrix (tile) and a second source matrix (tile), and stores the results into a destination matrix (tile). The size of the data elements to be loaded can vary depending on the instruction and tile support. Exemplary sizes include, but are not limited to: 16-bit, 32-bit, 64-bit, 128-bit, 256-bit, etc. The TDP4BIT instruction, however, operates on nibbles within quadword elements. In some embodiments, elements of rows of the destination matrix (tile) that do not have corresponding columns in the source matrix (tile) are zeroed.

## I. Exemplary Execution

[0165] FIG. 21 illustrates an exemplary execution of a TDP4BIT instruction. The TDP4BIT instruction format includes fields for an opcode and first source, second source, and destination identifiers, each to identify a matrix (tile) having a PAIR parameter set to FALSE. The identified destination matrix (tile) has M rows and N columns of elements. The identified first source matrix (tile) has M rows and K columns, and the identified second source matrix (tile) has K rows and N columns. As shown, a decoded TDP4BIT instruction **2102** is received by execution circuitry **2108**, which, in some embodiments, uses a grid of FMAs **2110** to perform a flow K times for each element (M,N) of the identified destination matrix **2118**, the flow including generating eight products by using multipliers **2112A-H** to multiply each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix; and using accumulator **2114** to accumulate and using saturation circuit **2116** to saturate the eight products with previous contents of the doubleword element (M,N) of the identified destination matrix.

[0166] In some embodiments, multipliers **2112A-H** perform all of the multiplications in parallel. In some embodiments, multipliers **2112A-H** perform all of the multiplications serially, over one or more additional cycles.

[0167] In some embodiments, the execution circuitry **2108** regularly saves state so that the execution circuitry **2108**, in case of a fault occurring in the midst of the operation, can resume execution after the fault is resolved and pick up where it left off. In such embodiments, the saved state includes one or more of a row pointer, a chunk pointer, a byte pointer, and an element pointer to identify the current destination tile element being loaded.

[0168] As detailed earlier, the source and destination matrices {tiles} may be loaded in a collection of registers, locations in memory, or in other storage accessible to execution circuitry. As shown, however, the identified source and destination matrices (tiles) are stored in a collection of registers. [0169] As shown, instruction **2102** includes a TDP4BIT opcode, a destination tile identifier, “tdest,” and first and second source tile identifiers, “tsrc1” and “tsrc2.”

[0170] In some embodiments, a matrix (tile) is configured to use only a subset of the rows and columns possible. For example, a matrix (tile) may have up to 16 rows and columns to use, but only use 4 of each. The configuration of each matrix (tile) is typically done by the execution of a configuration instruction prior to matrix (tile) usage.

## II. Exemplary Instruction Format(s)

[0171] An embodiment of a format for a TDP4BIT instruction is TDP4BIT tdest, tsrc1, tsrc2. In some embodiments, TDP4BIT is the opcode mnemonic of the instruction where “4BIT” indicates that 4-bit data nibbles are operated on. In some embodiments, the tdest field is a R/M value (such as **2546** of FIGS. **25A-B**), the tsrc1 field is REG **2544** of FIGS. **25A-B**. In some embodiments, the TDP4BIT opcode includes indicators, such as [U, S][U, S] prefixes or suffixes, to indicate whether each of the first and second identified sources is signed or unsigned.

[0172] In some embodiments, encodings of the instruction include a scale-index-base (SIB) type memory addressing operand that indirectly identifies multiple indexed destination locations in memory (e.g., field **2550** of FIGS. **25A-B**). In one embodiment, an SIB type memory operand may include an encoding identifying a base address register. The contents of the base address register may represent a base address in memory from which the addresses of the particular destination locations in memory are calculated. For example, the base address may be the address of the first location in a block of potential destination locations for an extended vector instruction. In one embodiment, a SIB type memory operand may include an encoding identifying an index register. Each element of the index register may specify an index or offset value usable to compute, from the base address, an address of a respective destination location within a block of potential destination locations. In one embodiment, a SIB type memory operand may include an encoding specifying a scaling factor to be applied to each index value when computing a respective destination address. For example, if a scaling factor value of four is encoded in the SIB type memory operand, each index value obtained from an element of the index register may be multiplied by four and then added to the base address to compute a destination address.

[0173] In one embodiment, an SIB type memory operand of the form vm32{x, y, z} may identify a vector array of memory operands specified using SIB type memory addressing. In this example, the array of memory addresses is specified using a common base register, a constant scaling factor, and a vector index register containing individual elements, each of which is a 32-bit index value. The vector index register may be a 128-bit register (e.g., XMM) register (vm32x), a 256-bit (e.g., YMM) register (vm32y), or a 512-bit (e.g., ZMM) register (vm32z). In another embodiment, an SIB type memory operand of the form vm64{x, y, z} may identify a vector array of memory operands specified using SIB type memory addressing. In this example, the array of memory addresses is specified using a common base register, a constant scaling factor, and a vector index register containing individual elements, each of which is a 64-bit index value. The vector index register may be a 128-bit register (e.g., XMM) register (vm64x), a 256-bit (e.g., YMM) register (vm64y) or a 512-bit (e.g., ZMM) register (vm64z).

## III. Exemplary Method(s) of Execution

[0174] FIG. **22** illustrates an embodiment of a processor executing a TDP4BIT instruction. At **2201**, an instruction is fetched. For example, a TDP4BIT instruction is fetched having fields for an opcode, a destination identifier to identify a M by N destination matrix, a first source identifier to identify a M by K first source matrix, and a second source identifier to identify a K by N second source matrix, where each of the matrices contains doubleword elements. In some embodiments, the instruction is fetched by fetch circuitry from an instruction cache. The opcode of the TDP4BIT

instruction indicates performing a packed dot product of element positions of the first and second identified source matrices (tiles). The source matrices contain doubleword elements, each of which is partitioned into 8 logical 4-bit nibbles.

[0175] The fetched instruction is decoded at **2203**. For example, the fetched TDP4BIT instruction is decoded by decode circuitry such as that detailed herein.

[0176] Execution of the decoded instruction is scheduled (as needed) at **2205**, which is an optional step (as indicated by its dashed border) insofar as it may occur at a different time, or not at all.

[0177] At **2207**, the decoded TDP4BIT instruction is executed by execution circuitry (hardware) such as that detailed herein. For the TDP4BIT instruction, the execution will cause execution circuitry the decoded instruction by performing a flow K times for each element (M,N) of the identified destination matrix, the flow comprising generating eight products by multiplying each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix; and accumulating and saturating the eight products with previous contents of the doubleword element (M,N) of the identified destination matrix. In some embodiments, the instruction is committed or retired at **2209**, which is an optional step (as indicated by its dashed border) insofar as it may occur at a different time, or not at all.

[0178] FIG. **23** illustrates a more detailed description of an execution of a TDP4BIT instruction. Typically, this is performed by execution circuitry such as that detailed above.

[0179] At **2302**, a TDP4BIT instruction is fetched having fields for an opcode and first source, second source, and destination identifiers, each to identify a matrix (tile) having a PAIR parameter set to FALSE.

[0180] At **2304**, a determination of whether ALL of the following is true is made whether any one or more errors are detected, in which case, a fault is generated at **2306**.

[0181] When all of the errors checks tested at **2304** pass, execution circuitry at **2308** For each element (M,N) of the identified destination matrix executes a flow K times to execute the remaining blocks of the flow. At **2310**, the execution circuitry generates eight products by multiplying each nibble of a doubleword element (M, K) of the identified first source matrix by a corresponding nibble of a doubleword element (K, N) of the identified second source matrix, and accumulates and saturates the eight products with previous contents of the doubleword element (M, N) of the identified destination matrix. At **2312**, N is incremented and, if any more columns remain, execution returns to **2308** to process the next column. Otherwise, at **2314**, K is incremented and, if any more K remain, execution returns to **2308** to process the next K. Otherwise, at **2316**, M is incremented and, if any more rows remain, execution returns to **2308** to process the next M. Otherwise, the process ends.

#### IV. Exemplary Pseudocode

[0182] FIG. **24A** is exemplary pseudocode describing an embodiment of a processor executing a TDP4BIT instruction. As shown in pseudocode **2400**, the TDP4BIT instruction includes an opcode, TDP4BIT, a destination matrix (tile) identifier to identify a M by N destination matrix (tile), “tdest,” a first source matrix (tile) identifier, “tsrc1,” to identify a M by K first source matrix (tile), and a second source matrix (tile) identifiers, “tsrc2,” to identify a K by N second source matrix (tile).

[0183] In some embodiments, each of the identified first source, second source, and destination matrices uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0184] As shown, the pseudocode **2400** first causes the execution circuitry to generate a fault if any of several error checks fails. For example, a fault is generated if either of tdest, tsrc1, and tsrc2 has a PAIR parameter set to TRUE.

[0185] Pseudocode **2400** further defines an “extend\_src” macro to sign-extend a signed source data value or to zero-extend an unsigned source data value. Pseudocode **2400** further defines a

“saturation\_fn” macro to perform unsigned saturation if both source data values are unsigned, and to perform signed saturation if either or both of the source data values are signed. For signed saturation, when an individual result is beyond the range of a signed doubleword integer (that is, greater than 7FFF\_FFFFH or less than 8000\_0000H), the saturated value of 7FFF\_FFFFH or 8000\_0000H, respectively, is written to the destination operand. For unsigned saturation, when an individual result value is beyond the range of an unsigned doubleword (that is, greater than FFFF\_FFFFH), the saturated unsigned doubleword integer value of FFFF\_FFFFH stored in the doubleword destination.

[0186] Pseudocode **2400**, when executed by a processor, causes the processor to compute the dot product of two doubleword operands, each containing eight 4-bit (nibbles) elements, with doubleword accumulation and saturation. As explained above, the 4-bit elements are signed or unsigned, and the saturation is signed if either operand is signed. Pseudocode **2400** causes the processor to perform a number of nested for loops to performing a flow K times for each element (M, N) of the identified destination matrix, the flow including generating eight products by multiplying each sign-extended or zero-extended nibble of a doubleword element (M, K) of the identified first source matrix by a corresponding sign-extended or zero-extended nibble of a doubleword element (K, N) of the identified second source matrix. The result of each nibble multiplication is stored in a doubleword, tprod[b], so will maintain exact precision. Pseudocode **2400** further causes the processor to accumulate and saturate (signed or unsigned) the eight products with previous contents of the doubleword element (M, N) of the identified destination matrix. Pseudocode **2400** causes the processor to zero any remaining configured rows of the destination matrix (tile), and reset pointers for a next invocation of the function.

[0187] In some embodiments, the execution circuitry saves a state after performing the flow on each element (M, N) of the identified destination matrix, and, in the case of a fault, uses the saved state after recovering from the fault to continue execution, wherein the state comprises at least a destination matrix row pointer and a destination matrix element pointer.

[0188] FIG. **24B** is exemplary pseudocode describing an embodiment of helper functions used by a processor executing a TDP4BIT instruction. Pseudocode **2450** defines the “write\_row\_and\_zero (treg, r, data, nbytes)” function. Pseudocode **2452** defines the “zero\_upper\_rows (treg, r)” function. Pseudocode **2454** defines the “zero\_tileconfig\_start ( )” function.

[0189] Pseudocode in FIG. **24A** and FIG. **24B** is self-documenting by virtue of its included function names and variable names.

#### Further Examples

[0190] Example 1 provides a processor including: decode circuitry to decode a tile dot product instruction having fields for an opcode, a destination identifier to identify a M by N destination matrix, a first source identifier to identify a M by K first source matrix, and a second source identifier to identify a K by N second source matrix, each of the matrices to contain doubleword elements, and execution circuitry to execute the decoded instruction by performing a flow K times for each element (M,N) of the identified destination matrix, the flow including: generating eight products by multiplying each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix; and accumulating and saturating the eight products with previous contents of the doubleword element (M,N) of the identified destination matrix

[0191] Example 2 includes the substance of the exemplary processor of Example 1, wherein the execution circuitry is further to generate a fault when a fault condition occurs, the fault condition including any one or more of: one or more of the first and second identified source matrices having a PAIR parameter set to TRUE, one or more of the first and second identified source matrices having a VALID parameter not set to TRUE, the identified destination matrix has a different number of rows than a number of rows of the identified first source matrix, the identified destination matrix has a different number of columns than a number of columns of the identified

second source matrix, and one or more of the identified first source, second source, and destination matrices have dimensions that exceed maximum dimensions, the dimensions including numbers of matrix rows and columns.

[0192] Example 3 includes the substance of the exemplary processor of Example 1, wherein the execution circuitry is further to either sign-extend or zero-extend each nibble being multiplied, based on the sign of the nibble.

[0193] Example 4 includes the substance of the exemplary processor of Example 1, wherein the saturating includes either signed saturation or unsigned saturation, depending on the sign of the eight products.

[0194] Example 5 includes the substance of the exemplary processor of Example 1, wherein a result of each multiplication of nibbles is stored in a doubleword register.

[0195] Example 6 includes the substance of the exemplary processor of any of Examples 1-5, wherein M is one of 2, 4, 8, and 16, N is one of 2, 4, 8, 16, and 32, and K is one of 1, 2, 4, 8, and 16.

[0196] Example 7 includes the substance of the exemplary processor of any of Examples 1-5, wherein the execution circuitry saves a state after performing the flow on each element (M, N) of the identified destination matrix, and, in the case of a fault, uses the saved state after recovering from the fault to continue execution of the K flows; wherein the state includes at least a destination matrix row pointer and a destination matrix element pointer.

[0197] Example 8 includes the substance of the exemplary processor of any of Examples 1-5, wherein the identified first and second source matrices each uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0198] Example 9 includes the substance of the exemplary processor of any of Examples 1-5, wherein the identified destination matrix uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0199] Example 10 includes the substance of the exemplary processor of any of Examples 1-5, wherein the opcode further indicates whether each of the first and second source matrices contains signed or unsigned elements.

[0200] Example 11 provides a method including: decoding, by decode circuitry, a tile dot product instruction having fields for an opcode, a destination identifier to identify a M by N destination matrix, a first source identifier to identify a M by K first source matrix, and a second source identifier to identify a K by N second source matrix, each of the matrices to contain doubleword elements, and executing, by execution circuitry, the decoded instruction by performing a flow K times for each element (M,N) of the identified destination matrix, the flow including: generating eight products by multiplying each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix; and accumulating and saturating the eight products with previous contents of the doubleword element (M,N) of the identified destination matrix

[0201] Example 12 includes the substance of the exemplary method of Example 11, further including generating a fault, by the execution circuitry, when a fault condition occurs, the fault condition including any one or more of: one or more of the first and second identified source matrices having a PAIR parameter set to TRUE, one or more of the first and second identified source matrices having a VALID parameter not set to TRUE, the identified destination matrix has a different number of rows than a number of rows of the identified first source matrix, the identified destination matrix has a different number of columns than a number of columns of the identified second source matrix, and one or more of the identified first source, second source, and destination matrices have dimensions that exceed maximum dimensions, the dimensions including numbers of matrix rows and columns.

[0202] Example 13 includes the substance of the exemplary method of Example 11, further including either sign-extending or zero-extending, by the execution circuitry, each nibble being

multipplied, based on the sign of the nibble.

[0203] Example 14 includes the substance of the exemplary method of Example 11, wherein the saturating includes either signed saturation or unsigned saturation, depending on the sign of the eight products.

[0204] Example 15 includes the substance of the exemplary method of Example 11, further including storing a result of each multiplication, by the execution circuitry, of nibbles in a doubleword register.

[0205] Example 16 includes the substance of the exemplary method of any of Examples 11-15, wherein M is one of 2, 4, 8, and 16, N is one of 2, 4, 8, 16, and 32, and K is one of 1, 2, 4, 8, and 16.

[0206] Example 17 includes the substance of the exemplary method of any of Examples 11-15, further including saving a state after performing the flow on each element (M, N) of the identified destination matrix, and, in the case of a fault, uses the saved state after recovering from the fault to continue execution of the K flows; wherein the state includes at least a destination matrix row pointer and a destination matrix element pointer.

[0207] Example 18 includes the substance of the exemplary method of any of Examples 11-15, wherein the identified first and second source matrices each uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0208] Example 19 includes the substance of the exemplary method of any of Examples 11-15, wherein the identified destination matrix uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0209] Example 20 includes the substance of the exemplary method of any of Examples 11-15, wherein the opcode further indicates whether each of the first and second source matrices contains signed or unsigned elements.

[0210] Example 21 provides a system including: a memory, and a processor including: decode circuitry to decode a tile dot product instruction having fields for an opcode, a destination identifier to identify a M by N destination matrix, a first source identifier to identify a M by K first source matrix, and a second source identifier to identify a K by N second source matrix, each of the matrices to contain doubleword elements, and execution circuitry to execute the decoded instruction by performing a flow K times for each element (M,N) of the identified destination matrix, the flow including: generating eight products by multiplying each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix; and accumulating and saturating the eight products with previous contents of the doubleword element (M,N) of the identified destination matrix

[0211] Example 22 includes the substance of the exemplary system of Example 21, wherein the execution circuitry is further to generate a fault when a fault condition occurs, the fault condition including any one or more of: one or more of the first and second identified source matrices having a PAIR parameter set to TRUE, one or more of the first and second identified source matrices having a VALID parameter not set to TRUE, the identified destination matrix has a different number of rows than a number of rows of the identified first source matrix, the identified destination matrix has a different number of columns than a number of columns of the identified second source matrix, and one or more of the identified first source, second source, and destination matrices have dimensions that exceed maximum dimensions, the dimensions including numbers of matrix rows and columns.

[0212] Example 23 includes the substance of the exemplary system of Example 21, wherein the execution circuitry is further to either sign-extend or zero-extend each nibble being multiplied, based on the sign of the nibble.

[0213] Example 24 includes the substance of the exemplary system of Example 21, wherein the saturating includes either signed saturation or unsigned saturation, depending on the sign of the

eight products.

[0214] Example 25 includes the substance of the exemplary system of Example 21, wherein a result of each multiplication of nibbles is stored in a doubleword register.

[0215] Example 26 includes the substance of the exemplary system of any of Examples 21-25, wherein M is one of 2, 4, 8, and 16, N is one of 2, 4, 8, 16, and 32, and K is one of 1, 2, 4, 8, and 16.

[0216] Example 27 includes the substance of the exemplary system of any of Examples 21-25, wherein the execution circuitry saves a state after performing the flow on each element (M, N) of the identified destination matrix, and, in the case of a fault, uses the saved state after recovering from the fault to continue execution of the K flows; wherein the state includes at least a destination matrix row pointer and a destination matrix element pointer.

[0217] Example 28 includes the substance of the exemplary system of any of Examples 21-25, wherein the identified first and second source matrices each uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0218] Example 29 includes the substance of the exemplary system of any of Examples 21-25, wherein the identified destination matrix uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0219] Example 30 includes the substance of the exemplary processor of any of Examples 21-25, wherein the opcode further indicates whether each of the first and second source matrices contains signed or unsigned elements.

[0220] Example 31 provides a non-transitory machine-readable medium containing instructions that, when executed by a processor, cause the processor to: decode, by decode circuitry, a tile dot product instruction having fields for an opcode, a destination identifier to identify a M by N destination matrix, a first source identifier to identify a M by K first source matrix, and a second source identifier to identify a K by N second source matrix, each of the matrices to contain doubleword elements, and execute, by execution circuitry, the decoded instruction by performing a flow K times for each element (M,N) of the identified destination matrix, the flow including: generating eight products by multiplying each nibble of a doubleword element (M,K) of the identified first source matrix by a corresponding nibble of a doubleword element (K,N) of the identified second source matrix; and accumulating and saturating the eight products with previous contents of the doubleword element (M,N) of the identified destination matrix

[0221] Example 32 includes the substance of the exemplary non-transitory machine-readable medium of Example 31, further containing instructions that, when executed by the processor, cause the processor to generate a fault, by the execution circuitry, when a fault condition occurs, the fault condition including any one or more of: one or more of the first and second identified source matrices having a PAIR parameter set to TRUE, one or more of the first and second identified source matrices having a VALID parameter not set to TRUE, the identified destination matrix has a different number of rows than a number of rows of the identified first source matrix, the identified destination matrix has a different number of columns than a number of columns of the identified second source matrix, and one or more of the identified first source, second source, and destination matrices have dimensions that exceed maximum dimensions, the dimensions including numbers of matrix rows and columns.

[0222] Example 33 includes the substance of the exemplary non-transitory machine-readable medium of Example 31, further containing instructions that, when executed by the processor, cause the processor to either sign-extend or zero-extend, by the execution circuitry, each nibble being multiplied, based on the sign of the nibble.

[0223] Example 34 includes the substance of the exemplary non-transitory machine-readable medium of Example 31, wherein the saturating includes either signed saturation or unsigned saturation, depending on the sign of the eight products.

[0224] Example 35 includes the substance of the exemplary non-transitory machine-readable

medium of Example 31, further containing instructions that, when executed by the processor, cause the processor to store a result of each multiplication, by the execution circuitry, of nibbles in a doubleword register.

[0225] Example 36 includes the substance of the exemplary non-transitory machine-readable medium of any of Examples 31-35, wherein M is one of 2, 4, 8, and 16, N is one of 2, 4, 8, 16, and 32, and K is one of 1, 2, 4, 8, and 16.

[0226] Example 37 includes the substance of the exemplary non-transitory machine-readable medium of any of Examples 31-35, further containing instructions that, when executed by the processor, cause the processor to save a state after performing the flow on each element (M,N) of the identified destination matrix, and, in the case of a fault, to use the saved state after recovering from the fault to continue execution of the K flows; wherein the state includes at least a destination matrix row pointer and a destination matrix element pointer.

[0227] Example 38 includes the substance of the exemplary non-transitory machine-readable medium of any of Examples 31-35, wherein the identified first and second source matrices each uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0228] Example 39 includes the substance of the exemplary non-transitory machine-readable medium of any of Examples 31-35, wherein the identified destination matrix uses one of a collection of registers and a plurality of memory locations to represent a matrix.

[0229] Example 40 includes the substance of the exemplary non-transitory machine-readable medium of any of Examples 31-35, wherein the opcode further indicates whether each of the first and second source matrices contains signed or unsigned elements.

## V. Detailed Exemplary Systems, Processors, and Emulation

[0230] Detailed herein are examples of hardware, software, etc. to execute the above described instructions. For example, what is described below details aspects of instruction execution including various pipeline stages such as fetch, decode, schedule, execute, retire, etc.

### Instruction Sets

[0231] An instruction set may include one or more instruction formats. A given instruction format may define various fields (e.g., number of bits, location of bits) to specify, among other things, the operation to be performed (e.g., opcode) and the operand(s) on which that operation is to be performed and/or other data field(s) (e.g., mask). Some instruction formats are further broken down though the definition of instruction templates (or subformats). For example, the instruction templates of a given instruction format may be defined to have different subsets of the instruction format's fields (the included fields are typically in the same order, but at least some have different bit positions because there are less fields included) and/or defined to have a given field interpreted differently. Thus, each instruction of an ISA is expressed using a given instruction format (and, if defined, in a given one of the instruction templates of that instruction format) and includes fields for specifying the operation and the operands. For example, an exemplary ADD instruction has a specific opcode and an instruction format that includes an opcode field to specify that opcode and operand fields to select operands (source1/destination and source2); and an occurrence of this ADD instruction in an instruction stream will have specific contents in the operand fields that select specific operands. A set of SIMD extensions referred to as the Advanced Vector Extensions (AVX) (AVX1 and AVX2) and using the Vector Extensions (VEX) coding scheme has been released and/or published (e.g., see Intel® 64 and IA-32 Architectures Software Developer's Manual, September 2014; and see Intel® Advanced Vector Extensions Programming Reference, October 2014).

### Exemplary Instruction Formats

[0232] Embodiments of the instruction(s) described herein may be embodied in different formats. Additionally, exemplary systems, architectures, and pipelines are detailed below. Embodiments of the instruction(s) may be executed on such systems, architectures, and pipelines, but are not limited to those detailed.



#### Generic Vector Friendly Instruction Format

[0233] A vector friendly instruction format is an instruction format that is suited for vector instructions (e.g., there are certain fields specific to vector operations). While embodiments are described in which both vector and scalar operations are supported through the vector friendly instruction format, alternative embodiments use only vector operations the vector friendly instruction format.

[0234] FIGS. 25A-25B are block diagrams illustrating a generic vector friendly instruction format and instruction templates thereof according to embodiments of the invention. FIG. 25A is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to embodiments of the invention; while FIG. 25B is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to embodiments of the invention. Specifically, a generic vector friendly instruction format 2500 for which are defined class A and class B instruction templates, both of which include no memory access 2505 instruction templates and memory access 2520 instruction templates. The term generic in the context of the vector friendly instruction format refers to the instruction format not being tied to any specific instruction set.

[0235] While embodiments of the invention will be described in which the vector friendly instruction format supports the following: a 64 byte vector operand length (or size) with 32 bit (4 byte) or 64 bit (8 byte) data element widths (or sizes) (and thus, a 64 byte vector consists of either 16 doubleword-size elements or alternatively, 8 quadword-size elements); a 64 byte vector operand length (or size) with 16 bit (2 byte) or 8 bit (1 byte) data element widths (or sizes); a 32 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); and a 16 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); alternative embodiments may support more, less and/or different vector operand sizes (e.g., 256 byte vector operands) with more, less, or different data element widths (e.g., 128 bit (16 byte) data element widths).

[0236] The class A instruction templates in FIG. 25A include: 1) within the no memory access 2505 instruction templates there is shown a no memory access, full round control type operation 2510 instruction template and a no memory access, data transform type operation 2515 instruction template; and 2) within the memory access 2520 instruction templates there is shown a memory access, temporal 2525 instruction template and a memory access, non-temporal 2530 instruction template. The class B instruction templates in FIG. 25B include: 1) within the no memory access 2505 instruction templates there is shown a no memory access, write mask control, partial round control type operation 2512 instruction template and a no memory access, write mask control, vsize type operation 2517 instruction template; and 2) within the memory access 2520 instruction templates there is shown a memory access, write mask control 2527 instruction template.

[0237] The generic vector friendly instruction format 2500 includes the following fields listed below in the order illustrated in FIGS. 25A-25B.

[0238] Format field 2540—a specific value (an instruction format identifier value) in this field uniquely identifies the vector friendly instruction format, and thus occurrences of instructions in the vector friendly instruction format in instruction streams. As such, this field is optional in the sense that it is not needed for an instruction set that has only the generic vector friendly instruction format.

[0239] Base operation field 2542—its content distinguishes different base operations.

[0240] Register index field 2544—its content, directly or through address generation, specifies the locations of the source and destination operands, be they in registers or in memory. These include a sufficient number of bits to select N registers from a PxQ (e.g. 32×512, 16×128, 32×1024, 64×1024) register file. While in one embodiment N may be up to three sources and one destination register, alternative embodiments may support more or less sources and destination registers (e.g., may support up to two sources where one of these sources also acts as the destination, may support

up to three sources where one of these sources also acts as the destination, may support up to two sources and one destination).

[0241] Modifier field **2546**—its content distinguishes occurrences of instructions in the generic vector instruction format that specify memory access from those that do not; that is, between no memory access **2505** instruction templates and memory access **2520** instruction templates. Memory access operations read and/or write to the memory hierarchy (in some cases specifying the source and/or destination addresses using values in registers), while non-memory access operations do not (e.g., the source and destinations are registers). While in one embodiment this field also selects between three different ways to perform memory address calculations, alternative embodiments may support more, less, or different ways to perform memory address calculations.

[0242] Augmentation operation field **2550**—its content distinguishes which one of a variety of different operations to be performed in addition to the base operation. This field is context specific. In one embodiment of the invention, this field is divided into a class field **2568**, an alpha field **2552**, and a beta field **2554**. The augmentation operation field **2550** allows common groups of operations to be performed in a single instruction rather than 2, 3, or 4 instructions.

[0243] Scale field **2560**—its content allows for the scaling of the index field's content for memory address generation (e.g., for address generation that uses  $2.\text{sup.scale} \times \text{index} + \text{base}$ ).

[0244] Displacement Field **2562A**—its content is used as part of memory address generation (e.g., for address generation that uses  $2.\text{sup.scale} \times \text{index} + \text{base} + \text{displacement}$ ).

[0245] Displacement Factor Field **2562B** (note that the juxtaposition of displacement field **2562A** directly over displacement factor field **2562B** indicates one or the other is used)—its content is used as part of address generation; it specifies a displacement factor that is to be scaled by the size of a memory access (N)—where N is the number of bytes in the memory access (e.g., for address generation that uses  $2.\text{sup.scale} \times \text{index} + \text{base} + \text{scaled displacement}$ ). Redundant low-order bits are ignored and hence, the displacement factor field's content is multiplied by the memory operands total size (N) in order to generate the final displacement to be used in calculating an effective address. The value of N is determined by the processor hardware at runtime based on the full opcode field **2574** (described later herein) and the data manipulation field **2554C**. The displacement field **2562A** and the displacement factor field **2562B** are optional in the sense that they are not used for the no memory access **2505** instruction templates and/or different embodiments may implement only one or none of the two.

[0246] Data element width field **2564**—its content distinguishes which one of a number of data element widths is to be used (in some embodiments for all instructions; in other embodiments for only some of the instructions). This field is optional in the sense that it is not needed if only one data element width is supported and/or data element widths are supported using some aspect of the opcodes.

[0247] Write mask field **2570**—its content controls, on a per data element position basis, whether that data element position in the destination vector operand reflects the result of the base operation and augmentation operation. Class A instruction templates support merging-writemasking, while class B instruction templates support both merging- and zeroing-writemasking. When merging, vector masks allow any set of elements in the destination to be protected from updates during the execution of any operation (specified by the base operation and the augmentation operation); in other one embodiment, preserving the old value of each element of the destination where the corresponding mask bit has a 0. In contrast, when zeroing vector masks allow any set of elements in the destination to be zeroed during the execution of any operation (specified by the base operation and the augmentation operation); in one embodiment, an element of the destination is set to 0 when the corresponding mask bit has a 0 value. A subset of this functionality is the ability to control the vector length of the operation being performed (that is, the span of elements being modified, from the first to the last one); however, it is not necessary that the elements that are modified be consecutive. Thus, the write mask field **2570** allows for partial vector operations,

including loads, stores, arithmetic, logical, etc. While embodiments of the invention are described in which the write mask field's **2570** content selects one of a number of write mask registers that contains the write mask to be used (and thus the write mask field's **2570** content indirectly identifies that masking to be performed), alternative embodiments instead or additionally allow the mask write field's **2570** content to directly specify the masking to be performed.

[0248] Immediate field **2572**—its content allows for the specification of an immediate. This field is optional in the sense that it is not present in an implementation of the generic vector friendly format that does not support immediate and it is not present in instructions that do not use an immediate.

[0249] Class field **2568**—its content distinguishes between different classes of instructions. With reference to FIGS. **25A-B**, the contents of this field select between class A and class B instructions. In FIGS. **25A-B**, rounded corner squares are used to indicate a specific value is present in a field (e.g., class A **2568A** and class B **2568B** for the class field **2568** respectively in FIGS. **25A-B**).

#### Instruction Templates of Class A

[0250] In the case of the non-memory access **2505** instruction templates of class A, the alpha field **2552** is interpreted as an RS field **2552A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **2552A.1** and data transform **2552A.2** are respectively specified for the no memory access, round type operation **2510** and the no memory access, data transform type operation **2515** instruction templates), while the beta field **2554** distinguishes which of the operations of the specified type is to be performed. In the no memory access **2505** instruction templates, the scale field **2560**, the displacement field **2562A**, and the displacement factor field **2562B** are not present.

#### No-Memory Access Instruction Templates—Full Round Control Type Operation

[0251] In the no memory access full round control type operation **2510** instruction template, the beta field **2554** is interpreted as a round control field **2554A**, whose content(s) provide static rounding. While in the described embodiments of the invention the round control field **2554A** includes a suppress all floating point exceptions (SAE) field **2556** and a round operation control field **2558**, alternative embodiments may support may encode both these concepts into the same field or only have one or the other of these concepts/fields (e.g., may have only the round operation control field **2558**).

[0252] SAE field **2556**—its content distinguishes whether or not to disable the exception event reporting; when the SAE field's **2556** content indicates suppression is enabled, a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler.

[0253] Round operation control field **2558**—its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **2558** allows for the changing of the rounding mode on a per instruction basis. In one embodiment of the invention where a processor includes a control register for specifying rounding modes, the round operation control field's **2550** content overrides that register value.

#### No Memory Access Instruction Templates—Data Transform Type Operation

[0254] In the no memory access data transform type operation **2515** instruction template, the beta field **2554** is interpreted as a data transform field **2554B**, whose content distinguishes which one of a number of data transforms is to be performed (e.g., no data transform, swizzle, broadcast).

[0255] In the case of a memory access **2520** instruction template of class A, the alpha field **2552** is interpreted as an eviction hint field **2552B**, whose content distinguishes which one of the eviction hints is to be used (in FIG. **25A**, temporal **2552B.1** and non-temporal **2552B.2** are respectively specified for the memory access, temporal **2525** instruction template and the memory access, non-temporal **2530** instruction template), while the beta field **2554** is interpreted as a data manipulation field **2554C**, whose content distinguishes which one of a number of data manipulation operations (also known as primitives) is to be performed (e.g., no manipulation; broadcast; up conversion of a

source; and down conversion of a destination). The memory access **2520** instruction templates include the scale field **2560**, and optionally the displacement field **2562A** or the displacement factor field **2562B**.

[0256] Vector memory instructions perform vector loads from and vector stores to memory, with conversion support. As with regular vector instructions, vector memory instructions transfer data from/to memory in a data element-wise fashion, with the elements that are actually transferred being dictated by the contents of the vector mask that is selected as the write mask.

#### Memory Access Instruction Templates—Temporal

[0257] Temporal data is data likely to be reused soon enough to benefit from caching. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

#### Memory Access Instruction Templates—Non-Temporal

[0258] Non-temporal data is data unlikely to be reused soon enough to benefit from caching in the 1st-level cache and should be given priority for eviction. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

#### Instruction Templates of Class B

[0259] In the case of the instruction templates of class B, the alpha field **2552** is interpreted as a write mask control (Z) field **2552C**, whose content distinguishes whether the write masking controlled by the write mask field **2570** should be a merging or a zeroing.

[0260] In the case of the non-memory access **2505** instruction templates of class B, part of the beta field **2554** is interpreted as an RL field **2557A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **2557A.1** and vector length (VSIZE) **2557A.2** are respectively specified for the no memory access, write mask control, partial round control type operation **2512** instruction template and the no memory access, write mask control, VSIZE type operation **2517** instruction template), while the rest of the beta field **2554** distinguishes which of the operations of the specified type is to be performed. In the no memory access **2505** instruction templates, the scale field **2560**, the displacement field **2562A**, and the displacement factor field **2562B** are not present.

[0261] In the no memory access, write mask control, partial round control type operation **2510** instruction template, the rest of the beta field **2554** is interpreted as a round operation field **2559A** and exception event reporting is disabled (a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler).

[0262] Round operation control field **2559A** —just as round operation control field **2558**, its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **2559A** allows for the changing of the rounding mode on a per instruction basis. In one embodiment of the invention where a processor includes a control register for specifying rounding modes, the round operation control field's **2550** content overrides that register value.

[0263] In the no memory access, write mask control, VSIZE type operation **2517** instruction template, the rest of the beta field **2554** is interpreted as a vector length field **2559B**, whose content distinguishes which one of a number of data vector lengths is to be performed on (e.g., 128, 256, or 512 byte).

[0264] In the case of a memory access **2520** instruction template of class B, part of the beta field **2554** is interpreted as a broadcast field **2557B**, whose content distinguishes whether or not the broadcast type data manipulation operation is to be performed, while the rest of the beta field **2554** is interpreted the vector length field **2559B**. The memory access **2520** instruction templates include the scale field **2560**, and optionally the displacement field **2562A** or the displacement factor field **2562B**.

[0265] With regard to the generic vector friendly instruction format **2500**, a full opcode field **2574** is shown including the format field **2540**, the base operation field **2542**, and the data element width

field **2564**. While one embodiment is shown where the full opcode field **2574** includes all of these fields, the full opcode field **2574** includes less than all of these fields in embodiments that do not support all of them. The full opcode field **2574** provides the operation code (opcode).

[0266] The augmentation operation field **2550**, the data element width field **2564**, and the write mask field **2570** allow these features to be specified on a per instruction basis in the generic vector friendly instruction format.

[0267] The combination of write mask field and data element width field create typed instructions in that they allow the mask to be applied based on different data element widths.

[0268] The various instruction templates found within class A and class B are beneficial in different situations. In some embodiments of the invention, different processors or different cores within a processor may support only class A, only class B, or both classes. For instance, a high performance general purpose out-of-order core intended for general-purpose computing may support only class B, a core intended primarily for graphics and/or scientific (throughput) computing may support only class A, and a core intended for both may support both (of course, a core that has some mix of templates and instructions from both classes but not all templates and instructions from both classes is within the purview of the invention). Also, a single processor may include multiple cores, all of which support the same class or in which different cores support different class. For instance, in a processor with separate graphics and general purpose cores, one of the graphics cores intended primarily for graphics and/or scientific computing may support only class A, while one or more of the general purpose cores may be high performance general purpose cores with out of order execution and register renaming intended for general-purpose computing that support only class B. Another processor that does not have a separate graphics core, may include one more general purpose in-order or out-of-order cores that support both class A and class B. Of course, features from one class may also be implemented in the other class in different embodiments of the invention. Programs written in a high level language would be put (e.g., just in time compiled or statically compiled) into a variety of different executable forms, including: 1) a form having only instructions of the class(es) supported by the target processor for execution; or 2) a form having alternative routines written using different combinations of the instructions of all classes and having control flow code that selects the routines to execute based on the instructions supported by the processor which is currently executing the code.

Exemplary Specific Vector Friendly Instruction Format

[0269] FIG. **26A** is a block diagram illustrating an exemplary specific vector friendly instruction format according to embodiments of the invention. FIG. **26A** shows a specific vector friendly instruction format **2600** that is specific in the sense that it specifies the location, size, interpretation, and order of the fields, as well as values for some of those fields. The specific vector friendly instruction format **2600** may be used to extend the x86 instruction set, and thus some of the fields are similar or the same as those used in the existing x86 instruction set and extension thereof (e.g., AVX). This format remains consistent with the prefix encoding field, real opcode byte field, MOD R/M field, SIB field, displacement field, and immediate fields of the existing x86 instruction set with extensions. The fields from FIG. **25** into which the fields from FIG. **26A** map are illustrated.

[0270] It should be understood that, although embodiments of the invention are described with reference to the specific vector friendly instruction format **2600** in the context of the generic vector friendly instruction format **2500** for illustrative purposes, the invention is not limited to the specific vector friendly instruction format **2600** except where claimed. For example, the generic vector friendly instruction format **2500** contemplates a variety of possible sizes for the various fields, while the specific vector friendly instruction format **2600** is shown as having fields of specific sizes. By way of specific example, while the data element width field **2564** is illustrated as a one bit field in the specific vector friendly instruction format **2600**, the invention is not so limited (that is, the generic vector friendly instruction format **2500** contemplates other sizes of the data element width field **2564**).

[0271] The generic vector friendly instruction format **2500** includes the following fields listed below in the order illustrated in FIG. **26A**.

[0272] EVEX Prefix (Bytes 0-3) **2602**—is encoded in a four-byte form.

[0273] Format Field **2540** (EVEX Byte 0, bits [7:0])—the first byte (EVEX Byte 0) is the format field **2540** and it contains 0x62 (the unique value used for distinguishing the vector friendly instruction format in one embodiment of the invention).

[0274] The second-fourth bytes (EVEX Bytes 1-3) include a number of bit fields providing specific capability.

[0275] REX field **2605** (EVEX Byte 1, bits [7-5])—consists of an EVEX.R bit field (EVEX Byte 1, bit [7]—R), EVEX.X bit field (EVEX byte 1, bit [6]—X), and EVEX.B bit field (EVEX byte 1, bit[5]—B). The EVEX.R, EVEX.X, and EVEX.B bit fields provide the same functionality as the corresponding VEX bit fields, and are encoded using 1s complement form, i.e. ZMM0 is encoded as 1111B, ZMM15 is encoded as 0000B. Other fields of the instructions encode the lower three bits of the register indexes as is known in the art (rrr, xxx, and bbb), so that Rrrr, Xxxx, and Bbbb may be formed by adding EVEX.R, EVEX.X, and EVEX.B.

[0276] REX' field **2610**—this is the first part of the REX' field and is the EVEX.R' bit field (EVEX Byte 1, bit [4]—R') that is used to encode either the upper 16 or lower 16 of the extended 32 register set. In one embodiment of the invention, this bit, along with others as indicated below, is stored in bit inverted format to distinguish (in the well-known x86 32-bit mode) from the BOUND instruction, whose real opcode byte is 62, but does not accept in the MOD R/M field (described below) the value of 11 in the MOD field; alternative embodiments of the invention do not store this and the other indicated bits below in the inverted format. A value of 1 is used to encode the lower 16 registers. In other words, R'Rrrr is formed by combining EVEX.R', EVEX.R, and the other RRR from other fields.

[0277] Opcode map field **2615** (EVEX byte 1, bits [3:0]—mmmm)—its content encodes an implied leading opcode byte (OF, OF **38**, or OF **3**).

[0278] Data element width field **2564** (EVEX byte 2, bit [7]—W)—is represented by the notation EVEX.W. EVEX.W is used to define the granularity (size) of the datatype (either 32-bit data elements or 64-bit data elements).

[0279] EVEX.vvvv **2620** (EVEX Byte 2, bits [6:3]—vvvv)—the role of EVEX.vvvv may include the following: 1) EVEX.vvvv encodes the first source register operand, specified in inverted (1s complement) form and is valid for instructions with 2 or more source operands; 2) EVEX.vvvv encodes the destination register operand, specified in 1s complement form for certain vector shifts; or 3) EVEX.vvvv does not encode any operand, the field is reserved and should contain 1111b. Thus, EVEX.vvvv field **2620** encodes the 4 low-order bits of the first source register specifier stored in inverted (1s complement) form. Depending on the instruction, an extra different EVEX bit field is used to extend the specifier size to 32 registers.

[0280] EVEX.U **2568** Class field (EVEX byte 2, bit [2]—U)—If EVEX.U=0, it indicates class A or EVEX.U0; if EVEX.U=1, it indicates class B or EVEX.U1.

[0281] Prefix encoding field **2625** (EVEX byte 2, bits [1:0]—pp)—provides additional bits for the base operation field. In addition to providing support for the legacy SSE instructions in the EVEX prefix format, this also has the benefit of compacting the SIMD prefix (rather than requiring a byte to express the SIMD prefix, the EVEX prefix requires only 2 bits). In one embodiment, to support legacy SSE instructions that use a SIMD prefix (66H, F2H, F3H) in both the legacy format and in the EVEX prefix format, these legacy SIMD prefixes are encoded into the SIMD prefix encoding field; and at runtime are expanded into the legacy SIMD prefix prior to being provided to the decoder's PLA (so the PLA can execute both the legacy and EVEX format of these legacy instructions without modification). Although newer instructions could use the EVEX prefix encoding field's content directly as an opcode extension, certain embodiments expand in a similar fashion for consistency but allow for different meanings to be specified by these legacy SIMD

prefixes. An alternative embodiment may redesign the PLA to support the 2 bit SIMD prefix encodings, and thus not require the expansion.

[0282] Alpha field **2552** (EVEX byte 3, bit [7]—EH; also known as EVEX.EH, EVEX.rs, EVEX.RL, EVEX.write mask control, and EVEX.N; also illustrated with  $\alpha$ )—as previously described, this field is context specific.

[0283] Beta field **2554** (EVEX byte 3, bits [6:4]-SSS, also known as EVEX.s.sub.2-0, EVEX.r.sub.2-0, EVEX.rr1, EVEX.LL0, EVEX.LLB; also illustrated with  $\beta\beta\beta$ )—as previously described, this field is context specific.

[0284] REX' field **2610**—this is the remainder of the REX' field and is the EVEX.V' bit field (EVEX Byte 3, bit [3]—V') that may be used to encode either the upper 16 or lower 16 of the extended 32 register set. This bit is stored in bit inverted format. A value of 1 is used to encode the lower 16 registers. In other words, V'VVVV is formed by combining EVEX.V', EVEX.vvvv.

[0285] Write mask field **2570** (EVEX byte 3, bits [2:0]—kkk)—its content specifies the index of a register in the write mask registers as previously described. In one embodiment of the invention, the specific value EVEX.kkk=000 has a special behavior implying no write mask is used for the particular instruction (this may be implemented in a variety of ways including the use of a write mask hardwired to all ones or hardware that bypasses the masking hardware).

[0286] Real Opcode Field **2630** (Byte 4) is also known as the opcode byte. Part of the opcode is specified in this field.

[0287] MOD R/M Field **2640** (Byte 5) includes MOD field **2642**, Reg field **2644**, and R/M field **2646**. As previously described, the MOD field's **2642** content distinguishes between memory access and non-memory access operations. The role of Reg field **2644** can be summarized to two situations: encoding either the destination register operand or a source register operand, or be treated as an opcode extension and not used to encode any instruction operand. The role of R/M field **2646** may include the following: encoding the instruction operand that references a memory address, or encoding either the destination register operand or a source register operand.

[0288] Scale, Index, Base (SIB) Byte (Byte 6)—As previously described, the scale field's **2550** content is used for memory address generation. SIB.xxx **2654** and SIB.bbb **2656**—the contents of these fields have been previously referred to with regard to the register indexes Xxxx and Bbbb.

[0289] Displacement field **2562A** (Bytes 7-10)—when MOD field **2642** contains 10, bytes 7-10 are the displacement field **2562A**, and it works the same as the legacy 32-bit displacement (disp32) and works at byte granularity.

[0290] Displacement factor field **2562B** (Byte 7)—when MOD field **2642** contains 01, byte 7 is the displacement factor field **2562B**. The location of this field is that same as that of the legacy x86 instruction set 8-bit displacement (disp8), which works at byte granularity. Since disp8 is sign extended, it can only address between -128 and 127 bytes offsets; in terms of 64 byte cache lines, disp8 uses 8 bits that can be set to only four really useful values -128, -64, 0, and 64; since a greater range is often needed, disp32 is used; however, disp32 requires 4 bytes. In contrast to disp8 and disp32, the displacement factor field **2562B** is a reinterpretation of disp8; when using displacement factor field **2562B**, the actual displacement is determined by the content of the displacement factor field multiplied by the size of the memory operand access (N). This type of displacement is referred to as disp8\*N. This reduces the average instruction length (a single byte of used for the displacement but with a much greater range). Such compressed displacement is based on the assumption that the effective displacement is multiple of the granularity of the memory access, and hence, the redundant low-order bits of the address offset do not need to be encoded. In other words, the displacement factor field **2562B** substitutes the legacy x86 instruction set 8-bit displacement. Thus, the displacement factor field **2562B** is encoded the same way as an x86 instruction set 8-bit displacement (so no changes in the ModRM/SIB encoding rules) with the only exception that disp8 is overloaded to disp8\*N. In other words, there are no changes in the encoding rules or encoding lengths but only in the interpretation of the displacement value by hardware

(which needs to scale the displacement by the size of the memory operand to obtain a byte-wise address offset). Immediate field **2572** operates as previously described.

#### Full Opcode Field

[0291] FIG. **26B** is a block diagram illustrating the fields of the specific vector friendly instruction format **2600** that make up the full opcode field **2574** according to one embodiment of the invention. Specifically, the full opcode field **2574** includes the format field **2540**, the base operation field **2542**, and the data element width (W) field **2564**. The base operation field **2542** includes the prefix encoding field **2625**, the opcode map field **2615**, and the real opcode field **2630**.

#### Register Index Field

[0292] FIG. **26C** is a block diagram illustrating the fields of the specific vector friendly instruction format **2600** that make up the register index field **2544** according to one embodiment of the invention. Specifically, the register index field **2544** includes the REX field **2605**, the REX' field **2610**, the MODR/M.reg field **2644**, the MODR/M.r/m field **2646**, the VVVV field **2620**, xxx field **2654**, and the bbb field **2656**.

#### Augmentation Operation Field

[0293] FIG. **26D** is a block diagram illustrating the fields of the specific vector friendly instruction format **2600** that make up the augmentation operation field **2550** according to one embodiment of the invention. When the class (U) field **2568** contains 0, it signifies EVEX.U0 (class A **2568A**); when it contains 1, it signifies EVEX.U1 (class B **2568B**). When U=0 and the MOD field **2642** contains 11 (signifying a no memory access operation), the alpha field **2552** (EVEX byte 3, bit [7]—EH) is interpreted as the rs field **2552A**. When the rs field **2552A** contains a 1 (round **2552A.1**), the beta field **2554** (EVEX byte 3, bits [6:4]—SSS) is interpreted as the round control field **2554A**. The round control field **2554A** includes a one bit SAE field **2556** and a two bit round operation field **2558**. When the rs field **2552A** contains a 0 (data transform **2552A.2**), the beta field **2554** (EVEX byte 3, bits [6:4]—SSS) is interpreted as a three bit data transform field **2554B**. When U=0 and the MOD field **2642** contains 00, 01, or 10 (signifying a memory access operation), the alpha field **2552** (EVEX byte 3, bit [7]—EH) is interpreted as the eviction hint (EH) field **2552B** and the beta field **2554** (EVEX byte 3, bits [6:4]—SSS) is interpreted as a three bit data manipulation field **2554C**.

[0294] When U=1, the alpha field **2552** (EVEX byte 3, bit [7]—EH) is interpreted as the write mask control (Z) field **2552C**. When U=1 and the MOD field **2642** contains 11 (signifying a no memory access operation), part of the beta field **2554** (EVEX byte 3, bit [4]—S.sub.0) is interpreted as the RL field **2557A**; when it contains a 1 (round **2557A.1**) the rest of the beta field **2554** (EVEX byte 3, bit [6-5]—S.sub.2-1) is interpreted as the round operation field **2559A**, while when the RL field **2557A** contains a 0 (VSIZE **2557A.2**) the rest of the beta field **2554** (EVEX byte 3, bit [6-5]—S.sub.2-1) is interpreted as the vector length field **2559B** (EVEX byte 3, bit [6-5]—L.sub.1-0). When U=1 and the MOD field **2642** contains 00, 01, or 10 (signifying a memory access operation), the beta field **2554** (EVEX byte 3, bits [6:4]—SSS) is interpreted as the vector length field **2559B** (EVEX byte 3, bit [6-5]—L.sub.1-0) and the broadcast field **2557B** (EVEX byte 3, bit [4]—B).

#### Exemplary Register Architecture

[0295] FIG. **27** is a block diagram of a register architecture **2700** according to one embodiment of the invention. In the embodiment illustrated, there are 32 vector registers **2710** that are 512 bits wide; these registers are referenced as zmm0 through zmm31. The lower order 256 bits of the lower 16 zmm registers are overlaid on registers ymm0-16. The lower order 128 bits of the lower 16 zmm registers (the lower order 128 bits of the ymm registers) are overlaid on registers xmm0-15. The specific vector friendly instruction format **2600** operates on these overlaid register file as illustrated in the below tables.

TABLE-US-00005 Adjustable Vector Length Class Operations Registers Instruction Templates A (FIG. 25A; 2510, 2515, zmm registers (the vector that do not include U = 0) 2525, 2530 length is



64 byte) the vector length B (FIG. 25B; 2512 zmm registers (the vector field 2559B U = 1) length is 64 byte) Instruction templates B (FIG. 25B; 2517, 2527 zmm, ymm, or xmm registers that do include U = 1) (the vector length is 64 the vector length byte, 32 byte, or 16 byte) field 2559B depending on the vector length field 2559B

[0296] In other words, the vector length field **2559B** selects between a maximum length and one or more other shorter lengths, where each such shorter length is half the length of the preceding length; and instructions templates without the vector length field **2559B** operate on the maximum vector length. Further, in one embodiment, the class B instruction templates of the specific vector friendly instruction format **2600** operate on packed or scalar single/double-precision floating point data and packed or scalar integer data. Scalar operations are operations performed on the lowest order data element position in a zmm/ymm/xmm register; the higher order data element positions are either left the same as they were prior to the instruction or zeroed depending on the embodiment.

[0297] Write mask registers **2715**—in the embodiment illustrated, there are 8 write mask registers (k0 through k7), each 64 bits in size. In an alternate embodiment, the write mask registers **2715** are 16 bits in size. As previously described, in one embodiment of the invention, the vector mask register k0 cannot be used as a write mask; when the encoding that would normally indicate k0 is used for a write mask, it selects a hardwired write mask of 0xFFFF, effectively disabling write masking for that instruction.

[0298] General-purpose registers **2725**—in the embodiment illustrated, there are sixteen 64-bit general-purpose registers that are used along with the existing x86 addressing modes to address memory operands. These registers are referenced by the names RAX, RBX, RCX, RDX, RBP, RSI, RDI, RSP, and R8 through R15.

[0299] Scalar floating point stack register file (x87 stack) **2745**, on which is aliased the MMX packed integer flat register file **2750**—in the embodiment illustrated, the x87 stack is an eight-element stack used to perform scalar floating-point operations on 32/64/80-bit floating point data using the x87 instruction set extension; while the MMX registers are used to perform operations on 64-bit packed integer data, as well as to hold operands for some operations performed between the MMX and XMM registers.

[0300] Alternative embodiments of the invention may use wider or narrower registers. Additionally, alternative embodiments of the invention may use more, less, or different register files and registers.

#### Exemplary Core Architectures, Processors, and Computer Architectures

[0301] Processor cores may be implemented in different ways, for different purposes, and in different processors. For instance, implementations of such cores may include: 1) a general purpose in-order core intended for general-purpose computing; 2) a high performance general purpose out-of-order core intended for general-purpose computing; 3) a special purpose core intended primarily for graphics and/or scientific (throughput) computing. Implementations of different processors may include: 1) a CPU including one or more general purpose in-order cores intended for general-purpose computing and/or one or more general purpose out-of-order cores intended for general-purpose computing; and 2) a coprocessor including one or more special purpose cores intended primarily for graphics and/or scientific (throughput). Such different processors lead to different computer system architectures, which may include: 1) the coprocessor on a separate chip from the CPU; 2) the coprocessor on a separate die in the same package as a CPU; 3) the coprocessor on the same die as a CPU (in which case, such a coprocessor is sometimes referred to as special purpose logic, such as integrated graphics and/or scientific (throughput) logic, or as special purpose cores); and 4) a system on a chip that may include on the same die the described CPU (sometimes referred to as the application core(s) or application processor(s)), the above described coprocessor, and additional functionality. Exemplary core architectures are described next, followed by descriptions of exemplary processors and computer architectures.

## Exemplary Core Architectures

### In-Order and Out-of-Order Core Block Diagram

[0302] FIG. **28A** is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to embodiments of the invention. FIG. **28B** is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to embodiments of the invention. The solid lined boxes in FIGS. **28A-B** illustrate the in-order pipeline and in-order core, while the optional addition of the dashed lined boxes illustrates the register renaming, out-of-order issue/execution pipeline and core. Given that the in-order aspect is a subset of the out-of-order aspect, the out-of-order aspect will be described.

[0303] In FIG. **28A**, a processor pipeline **2800** includes a fetch stage **2802**, a length decode stage **2804**, a decode stage **2806**, an allocation stage **2808**, a renaming stage **2810**, a scheduling (also known as a dispatch or issue) stage **2812**, a register read/memory read stage **2814**, an execute stage **2816**, a write back/memory write stage **2818**, an exception handling stage **2822**, and a commit stage **2824**.

[0304] FIG. **28B** shows processor core **2890** including a front end unit **2830** coupled to an execution engine unit **2850**, and both are coupled to a memory unit **2870**. The core **2890** may be a reduced instruction set computing (RISC) core, a complex instruction set computing (CISC) core, a very long instruction word (VLIW) core, or a hybrid or alternative core type. As yet another option, the core **2890** may be a special-purpose core, such as, for example, a network or communication core, compression engine, coprocessor core, general purpose computing graphics processing unit (GPGPU) core, graphics core, or the like.

[0305] The front end unit **2830** includes a branch prediction unit **2832** coupled to an instruction cache unit **2834**, which is coupled to an instruction translation lookaside buffer (TLB) **2836**, which is coupled to an instruction fetch unit **2838**, which is coupled to a decode unit **2840**. The decode unit **2840** (or decoder) may decode instructions, and generate as an output one or more micro-operations, micro-code entry points, microinstructions, other instructions, or other control signals, which are decoded from, or which otherwise reflect, or are derived from, the original instructions. The decode unit **2840** may be implemented using various different mechanisms. Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc. In one embodiment, the core **2890** includes a microcode ROM or other medium that stores microcode for certain macroinstructions (e.g., in decode unit **2840** or otherwise within the front end unit **2830**). The decode unit **2840** is coupled to a rename/allocator unit **2852** in the execution engine unit **2850**.

[0306] The execution engine unit **2850** includes the rename/allocator unit **2852** coupled to a retirement unit **2854** and a set of one or more scheduler unit(s) **2856**. The scheduler unit(s) **2856** represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) **2856** is coupled to the physical register file(s) unit(s) **2858**. Each of the physical register file(s) units **2858** represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point, status (e.g., an instruction pointer that is the address of the next instruction to be executed), etc. In one embodiment, the physical register file(s) unit **2858** comprises a vector registers unit, a write mask registers unit, and a scalar registers unit. These register units may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) unit(s) **2858** is overlapped by the retirement unit **2854** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement unit **2854** and the physical

register file(s) unit(s) **2858** are coupled to the execution cluster(s) **2860**. The execution cluster(s) **2860** includes a set of one or more execution units **2862** and a set of one or more memory access units **2864**. The execution units **2862** may perform various operations (e.g., shifts, addition, subtraction, multiplication) and on various types of data (e.g., scalar floating point, packed integer, packed floating point, vector integer, vector floating point). While some embodiments may include a number of execution units dedicated to specific functions or sets of functions, other embodiments may include only one execution unit or multiple execution units that all perform all functions. The scheduler unit(s) **2856**, physical register file(s) unit(s) **2858**, and execution cluster(s) **2860** are shown as being possibly plural because certain embodiments create separate pipelines for certain types of data/operations (e.g., a scalar integer pipeline, a scalar floating point/packed integer/packed floating point/vector integer/vector floating point pipeline, and/or a memory access pipeline that each have their own scheduler unit, physical register file(s) unit, and/or execution cluster—and in the case of a separate memory access pipeline, certain embodiments are implemented in which only the execution cluster of this pipeline has the memory access unit(s) **2864**). It should also be understood that where separate pipelines are used, one or more of these pipelines may be out-of-order issue/execution and the rest in-order.

[0307] The set of memory access units **2864** is coupled to the memory unit **2870**, which includes a data TLB unit **2872** coupled to a data cache unit **2874** coupled to a level 2 (L2) cache unit **2876**. In one exemplary embodiment, the memory access units **2864** may include a load unit, a store address unit, and a store data unit, each of which is coupled to the data TLB unit **2872** in the memory unit **2870**. The instruction cache unit **2834** is further coupled to a level 2 (L2) cache unit **2876** in the memory unit **2870**. The L2 cache unit **2876** is coupled to one or more other levels of cache and eventually to a main memory.

[0308] By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement the pipeline **2800** as follows: 1) the instruction fetch **2838** performs the fetch and length decoding stages **2802** and **2804**; 2) the decode unit **2840** performs the decode stage **2806**; 3) the rename/allocator unit **2852** performs the allocation stage **2808** and renaming stage **2810**; 4) the scheduler unit(s) **2856** performs the schedule stage **2812**; 5) the physical register file(s) unit(s) **2858** and the memory unit **2870** perform the register read/memory read stage **2814**; the execution cluster **2860** perform the execute stage **2816**; 6) the memory unit **2870** and the physical register file(s) unit(s) **2858** perform the write back/memory write stage **2818**; 7) various units may be involved in the exception handling stage **2822**; and 8) the retirement unit **2854** and the physical register file(s) unit(s) **2858** perform the commit stage **2824**.

[0309] The core **2890** may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, CA; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, CA), including the instruction(s) described herein. In one embodiment, the core **2890** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

[0310] It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyperthreading technology).

[0311] While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used in an in-order architecture. While the illustrated embodiment of the processor also includes separate instruction and data cache units **2834/2874** and a shared L2 cache unit **2876**, alternative embodiments may have a single internal cache for both

instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some embodiments, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache may be external to the core and/or the processor.

#### Specific Exemplary In-Order Core Architecture

[0312] FIGS. **29A-B** illustrate a block diagram of a more specific exemplary in-order core architecture, which core would be one of several logic blocks (including other cores of the same type and/or different types) in a chip. The logic blocks communicate through a high-bandwidth interconnect network (e.g., a ring network) with some fixed function logic, memory I/O interfaces, and other necessary I/O logic, depending on the application.

[0313] FIG. **29A** is a block diagram of a single processor core, along with its connection to the on-die interconnect network **2902** and with its local subset of the Level 2 (L2) cache **2904**, according to embodiments of the invention. In one embodiment, an instruction decoder **2900** supports the x86 instruction set with a packed data instruction set extension. An L1 cache **2906** allows low-latency accesses to cache memory into the scalar and vector units. While in one embodiment (to simplify the design), a scalar unit **2908** and a vector unit **2910** use separate register sets (respectively, scalar registers **2912** and vector registers **2914**) and data transferred between them is written to memory and then read back in from a level 1 (L1) cache **2906**, alternative embodiments of the invention may use a different approach (e.g., use a single register set or include a communication path that allow data to be transferred between the two register files without being written and read back).

[0314] The local subset of the L2 cache **2904** is part of a global L2 cache that is divided into separate local subsets, one per processor core. Each processor core has a direct access path to its own local subset of the L2 cache **2904**. Data read by a processor core is stored in its L2 cache subset **2904** and can be accessed quickly, in parallel with other processor cores accessing their own local L2 cache subsets. Data written by a processor core is stored in its own L2 cache subset **2904** and is flushed from other subsets, if necessary. The ring network ensures coherency for shared data. The ring network is bi-directional to allow agents such as processor cores, L2 caches and other logic blocks to communicate with each other within the chip. Each ring data-path is 1012-bits wide per direction.

[0315] FIG. **29B** is an expanded view of part of the processor core in FIG. **29A** according to embodiments of the invention. FIG. **29B** includes an L1 data cache **2906A** part of the L1 cache **2906**, as well as more detail regarding the vector unit **2910** and the vector registers **2914**.

Specifically, the vector unit **2910** is a 16-wide vector processing unit (VPU) (see the 16-wide ALU **2928**), which executes one or more of integer, single-precision float, and double-precision float instructions. The VPU supports swizzling the register inputs with swizzle unit **2920**, numeric conversion with numeric convert units **2922A-B**, and replication with replication unit **2924** on the memory input. Write mask registers **2926** allow predicating resulting vector writes.

[0316] FIG. **30** is a block diagram of a processor **3000** that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to embodiments of the invention. The solid lined boxes in FIG. **30** illustrate a processor **3000** with a single core **3002A**, a system agent **3010**, a set of one or more bus controller units **3016**, while the optional addition of the dashed lined boxes illustrates an alternative processor **3000** with multiple cores **3002A-N**, a set of one or more integrated memory controller unit(s) **3014** in the system agent unit **3010**, and special purpose logic **3008**.

[0317] Thus, different implementations of the processor **3000** may include: 1) a CPU with the special purpose logic **3008** being integrated graphics and/or scientific (throughput) logic (which may include one or more cores), and the cores **3002A-N** being one or more general purpose cores (e.g., general purpose in-order cores, general purpose out-of-order cores, a combination of the two); 2) a coprocessor with the cores **3002A-N** being a large number of special purpose cores intended primarily for graphics and/or scientific (throughput); and 3) a coprocessor with the cores **3002A-N**

being a large number of general purpose in-order cores. Thus, the processor **3000** may be a general-purpose processor, coprocessor or special-purpose processor, such as, for example, a network or communication processor, compression engine, graphics processor, GPGPU (general purpose graphics processing unit), a high-throughput many integrated core (MIC) coprocessor (including 30 or more cores), embedded processor, or the like. The processor **3000** may be a part of and/or may be implemented on one or more substrates using any of a number of process technologies, such as, for example, BiCMOS, CMOS, or NMOS.

[0318] The memory hierarchy includes one or more levels of cache within the cores, a set or one or more shared cache units **3006**, and external memory (not shown) coupled to the set of integrated memory controller units **3014**. The set of shared cache units **3006** may include one or more mid-level caches, such as level 2 (L2), level 3 (L3), level 4 (L4), or other levels of cache, a last level cache (LLC), and/or combinations thereof. While in one embodiment a ring based interconnect unit **3012** interconnects the integrated graphics logic **3008** (integrated graphics logic **3008** is an example of and is also referred to herein as special purpose logic), the set of shared cache units **3006**, and the system agent unit **3010**/integrated memory controller unit(s) **3014**, alternative embodiments may use any number of well-known techniques for interconnecting such units. In one embodiment, coherency is maintained between one or more cache units **3006** and cores **3002-A-N**.

[0319] In some embodiments, one or more of the cores **3002A-N** are capable of multi-threading. The system agent **3010** includes those components coordinating and operating cores **3002A-N**. The system agent unit **3010** may include for example a power control unit (PCU) and a display unit. The PCU may be or include logic and components needed for regulating the power state of the cores **3002A-N** and the integrated graphics logic **3008**. The display unit is for driving one or more externally connected displays.

[0320] The cores **3002A-N** may be homogenous or heterogeneous in terms of architecture instruction set; that is, two or more of the cores **3002A-N** may be capable of execution the same instruction set, while others may be capable of executing only a subset of that instruction set or a different instruction set.

#### Exemplary Computer Architectures

[0321] FIGS. **31-34** are block diagrams of exemplary computer architectures. Other system designs and configurations known in the arts for laptops, desktops, handheld PCs, personal digital assistants, engineering workstations, servers, network devices, network hubs, switches, embedded processors, digital signal processors (DSPs), graphics devices, video game devices, set-top boxes, micro controllers, cell phones, portable media players, hand held devices, and various other electronic devices, are also suitable. In general, a huge variety of systems or electronic devices capable of incorporating a processor and/or other execution logic as disclosed herein are generally suitable.

[0322] Referring now to FIG. **31**, shown is a block diagram of a system **3100** in accordance with one embodiment of the present invention. The system **3100** may include one or more processors **3110**, **3115**, which are coupled to a controller hub **3120**. In one embodiment the controller hub **3120** includes a graphics memory controller hub (GMCH) **3190** and an Input/Output Hub (IOH) **3150** (which may be on separate chips); the GMCH **3190** includes memory and graphics controllers to which are coupled memory **3140** and a coprocessor **3145**; the IOH **3150** couples input/output (I/O) devices **3160** to the GMCH **3190**. Alternatively, one or both of the memory and graphics controllers are integrated within the processor (as described herein), the memory **3140** and the coprocessor **3145** are coupled directly to the processor **3110**, and the controller hub **3120** in a single chip with the IOH **3150**.

[0323] The optional nature of additional processors **3115** is denoted in FIG. **31** with broken lines. Each processor **3110**, **3115** may include one or more of the processing cores described herein and may be some version of the processor **3000**.

[0324] The memory **3140** may be, for example, dynamic random access memory (DRAM), phase change memory (PCM), or a combination of the two. For at least one embodiment, the controller hub **3120** communicates with the processor(s) **3110**, **3115** via a multi-drop bus, such as a frontside bus (FSB), point-to-point interface such as QuickPath Interconnect (QPI), or similar connection **3195**.

[0325] In one embodiment, the coprocessor **3145** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like. In one embodiment, controller hub **3120** may include an integrated graphics accelerator.

[0326] There can be a variety of differences between the physical resources **3110**, **3115** in terms of a spectrum of metrics of merit including architectural, microarchitectural, thermal, power consumption characteristics, and the like.

[0327] In one embodiment, the processor **3110** executes instructions that control data processing operations of a general type. Embedded within the instructions may be coprocessor instructions. The processor **3110** recognizes these coprocessor instructions as being of a type that should be executed by the attached coprocessor **3145**. Accordingly, the processor **3110** issues these coprocessor instructions (or control signals representing coprocessor instructions) on a coprocessor bus or other interconnect, to coprocessor **3145**. Coprocessor(s) **3145** accept and execute the received coprocessor instructions.

[0328] Referring now to FIG. **32**, shown is a block diagram of a first more specific exemplary system **3200** in accordance with an embodiment of the present invention. As shown in FIG. **32**, multiprocessor system **3200** is a point-to-point interconnect system, and includes a first processor **3270** and a second processor **3280** coupled via a point-to-point interconnect **3250**. Each of processors **3270** and **3280** may be some version of the processor **3000**. In one embodiment of the invention, processors **3270** and **3280** are respectively processors **3110** and **3115**, while coprocessor **3238** is coprocessor **3145**. In another embodiment, processors **3270** and **3280** are respectively processor **3110** coprocessor **3145**.

[0329] Processors **3270** and **3280** are shown including integrated memory controller (IMC) units **3272** and **3282**, respectively. Processor **3270** also includes as part of its bus controller units point-to-point (P-P) interfaces **3276** and **3278**; similarly, second processor **3280** includes P-P interfaces **3286** and **3288**. Processors **3270**, **3280** may exchange information via a point-to-point (P-P) interface **3250** using P-P interface circuits **3278**, **3288**. As shown in FIG. **32**, IMCs **3272** and **3282** couple the processors to respective memories, namely a memory **3232** and a memory **3234**, which may be portions of main memory locally attached to the respective processors.

[0330] Processors **3270**, **3280** may each exchange information with a chipset **3290** via individual P-P interfaces **3252**, **3254** using point to point interface circuits **3276**, **3294**, **3286**, **3298**. Chipset **3290** may optionally exchange information with the coprocessor **3238** via a high-performance interface **3292**. In one embodiment, the coprocessor **3238** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like.

[0331] A shared cache (not shown) may be included in either processor or outside of both processors, yet connected with the processors via P-P interconnect, such that either or both processors' local cache information may be stored in the shared cache if a processor is placed into a low power mode.

[0332] Chipset **3290** may be coupled to a first bus **3216** via an interface **3296**. In one embodiment, first bus **3216** may be a Peripheral Component Interconnect (PCI) bus, or a bus such as a PCI Express bus or another third generation I/O interconnect bus, although the scope of the present invention is not so limited.

[0333] As shown in FIG. **32**, various I/O devices **3214** may be coupled to first bus **3216**, along with a bus bridge **3218** which couples first bus **3216** to a second bus **3220**. In one embodiment, one or

more additional processor(s) **3215**, such as coprocessors, high-throughput MIC processors, GPGPU's, accelerators (such as, e.g., graphics accelerators or digital signal processing (DSP) units), field programmable gate arrays, or any other processor, are coupled to first bus **3216**. In one embodiment, second bus **3220** may be a low pin count (LPC) bus. Various devices may be coupled to a second bus **3220** including, for example, a keyboard and/or mouse **3222**, communication devices **3227** and a storage unit **3228** such as a disk drive or other mass storage device which may include instructions/code and data **3230**, in one embodiment. Further, an audio I/O **3224** may be coupled to the second bus **3220**. Note that other architectures are possible. For example, instead of the point-to-point architecture of FIG. **32**, a system may implement a multi-drop bus or other such architecture.

[0334] Referring now to FIG. **33**, shown is a block diagram of a second more specific exemplary system **3300** in accordance with an embodiment of the present invention. Like elements in FIGS. **32** and **33** bear like reference numerals, and certain aspects of FIG. **32** have been omitted from FIG. **33** in order to avoid obscuring other aspects of FIG. **33**.

[0335] FIG. **33** illustrates that the processors **3270**, **3280** may include integrated memory and I/O control logic ("CL") **3272** and **3282**, respectively. Thus, the CL **3272**, **3282** include integrated memory controller units and include I/O control logic. FIG. **33** illustrates that not only are the memories **3232**, **3234** coupled to the CL **3272**, **3282**, but also that I/O devices **3314** are also coupled to the control logic **3272**, **3282**. Legacy I/O devices **3315** are coupled to the chipset **3290**.

[0336] Referring now to FIG. **34**, shown is a block diagram of a SoC **3400** in accordance with an embodiment of the present invention. Similar elements in FIG. **30** bear like reference numerals. Also, dashed lined boxes are optional features on more advanced SoCs. In FIG. **34**, an interconnect unit(s) **3402** is coupled to: an application processor **3410** which includes a set of one or more cores **3002A-N**, which include cache units **3004A-N**, and shared cache unit(s) **3006**; a system agent unit **3010**; a bus controller unit(s) **3016**; an integrated memory controller unit(s) **3014**; a set or one or more coprocessors **3420** which may include integrated graphics logic, an image processor, an audio processor, and a video processor; a static random access memory (SRAM) unit **3430**; a direct memory access (DMA) unit **3432**; and a display unit **3440** for coupling to one or more external displays. In one embodiment, the coprocessor(s) **3420** include a special-purpose processor, such as, for example, a network or communication processor, compression engine, GPGPU, a high-throughput MIC processor, embedded processor, or the like.

[0337] Embodiments of the mechanisms disclosed herein may be implemented in hardware, software, firmware, or a combination of such implementation approaches. Embodiments of the invention may be implemented as computer programs or program code executing on programmable systems comprising at least one processor, a storage system (including volatile and non-volatile memory and/or storage elements), at least one input device, and at least one output device.

[0338] Program code, such as code **3230** illustrated in FIG. **32**, may be applied to input instructions to perform the functions described herein and generate output information. The output information may be applied to one or more output devices, in known fashion. For purposes of this application, a processing system includes any system that has a processor, such as, for example; a digital signal processor (DSP), a microcontroller, an application specific integrated circuit (ASIC), or a microprocessor.

[0339] The program code may be implemented in a high level procedural or object oriented programming language to communicate with a processing system. The program code may also be implemented in assembly or machine language, if desired. In fact, the mechanisms described herein are not limited in scope to any particular programming language. In any case, the language may be a compiled or interpreted language.

[0340] One or more aspects of at least one embodiment may be implemented by representative instructions stored on a machine-readable medium which represents various logic within the processor, which when read by a machine causes the machine to fabricate logic to perform the

techniques described herein. Such representations, known as “IP cores,” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

[0341] Such machine-readable storage media may include, without limitation, non-transitory, tangible arrangements of articles manufactured or formed by a machine or device, including storage media such as hard disks, any other type of disk including floppy disks, optical disks, compact disk read-only memories (CD-ROMs), compact disk rewritables (CD-RWs), and magneto-optical disks, semiconductor devices such as read-only memories (ROMs), random access memories (RAMs) such as dynamic random access memories (DRAMs), static random access memories (SRAMs), erasable programmable read-only memories (EPROMs), flash memories, electrically erasable programmable read-only memories (EEPROMs), phase change memory (PCM), magnetic or optical cards, or any other type of media suitable for storing electronic instructions.

[0342] Accordingly, embodiments of the invention also include non-transitory, tangible machine-readable media containing instructions or containing design data, such as Hardware Description Language (HDL), which defines structures, circuits, apparatuses, processors and/or system features described herein. Such embodiments may also be referred to as program products.

Emulation (Including Binary Translation, Code Morphing, Etc.)

[0343] In some cases, an instruction converter may be used to convert an instruction from a source instruction set to a target instruction set. For example, the instruction converter may translate (e.g., using static binary translation, dynamic binary translation including dynamic compilation), morph, emulate, or otherwise convert an instruction to one or more other instructions to be processed by the core. The instruction converter may be implemented in software, hardware, firmware, or a combination thereof. The instruction converter may be on processor, off processor, or part on and part off processor.

[0344] FIG. 35 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the invention. In the illustrated embodiment, the instruction converter is a software instruction converter, although alternatively the instruction converter may be implemented in software, firmware, hardware, or various combinations thereof. FIG. 35 shows a program in a high level language 3502 may be compiled using an x86 compiler 3504 to generate x86 binary code 3506 that may be natively executed by a processor with at least one x86 instruction set core 3516. The processor with at least one x86 instruction set core 3516 represents any processor that can perform substantially the same functions as an Intel processor with at least one x86 instruction set core by compatibly executing or otherwise processing (1) a substantial portion of the instruction set of the Intel x86 instruction set core or (2) object code versions of applications or other software targeted to run on an Intel processor with at least one x86 instruction set core, in order to achieve substantially the same result as an Intel processor with at least one x86 instruction set core. The x86 compiler 3504 represents a compiler that is operable to generate x86 binary code 3506 (e.g., object code) that can, with or without additional linkage processing, be executed on the processor with at least one x86 instruction set core 3516. Similarly, FIG. 35 shows the program in the high level language 3502 may be compiled using an alternative instruction set compiler 3508 to generate alternative instruction set binary code 3510 that may be natively executed by a processor without at least one x86 instruction set core 3514 (e.g., a processor with cores that execute the MIPS instruction set of MIPS Technologies of Sunnyvale, CA and/or that execute the ARM instruction set of ARM Holdings of Sunnyvale, CA). The instruction converter 3512 is used to convert the x86 binary code 3506 into code that may be natively executed by the processor without an x86 instruction set core 3514. This converted code is not likely to be the same as the alternative instruction set binary code 3510 because an instruction converter capable of this is difficult to make; however, the converted code will accomplish the general operation and be



made up of instructions from the alternative instruction set. Thus, the instruction converter 3512 represents software, firmware, hardware, or a combination thereof that, through emulation, simulation or any other process, allows a processor or other electronic device that does not have an x86 instruction set processor or core to execute the x86 binary code 3506.

## Claims

**1.-20.** (canceled)

**21.** An apparatus comprising: a plurality of memory controllers; a level-two (L2) cache memory coupled to the plurality of memory controllers; and a processor coupled to the plurality of memory controllers, and coupled to the L2 cache memory, the processor having a plurality of cores comprising circuitry to perform operations corresponding to an instruction, the instruction to indicate a first matrix having M rows by K columns of 32-bit elements each having a plurality of 4-bit data elements, a second matrix having K rows by N columns of 32-bit elements each having a plurality of 4-bit data elements, and a third matrix having M rows by N columns, wherein the 4-bit data elements of the first and second matrices are floating-point data elements, the operations including to: for each data element position at a row m of the M rows and at a column n of the N columns of the third matrix: generate a dot product of the plurality of 4-bit data elements of each of K 32-bit elements of a row m of the first matrix and corresponding ones of the plurality of 4-bit data elements of each of K 32-bit elements of a column n of the second matrix; and store a 32-bit result data element at the data element position of the third matrix, the 32-bit result data element to be based on accumulation of the dot product with a 32-bit data element corresponding to the data element position.

**22.** The apparatus of claim 21, wherein the operations include to check whether the first matrix is one of a pair of matrices.

**23.** The apparatus of claim 21, wherein N is 8.

**24.** The apparatus of claim 22, wherein M is 16.

**25.** The apparatus of claim 21, wherein the first matrix is to be stored in a plurality of registers of the processor, and wherein the second matrix is to be stored in a plurality of registers of the processor.

**26.** The apparatus of claim 21, wherein the plurality of cores comprise graphics cores.

**27.** The apparatus of claim 21, wherein the plurality of cores comprise graphics cores, and wherein the processor comprises heterogeneous graphics cores.

**28.** The apparatus of claim 21, further comprising an interconnect interface coupled to the processor.

**29.** The apparatus of claim 21, further comprising an instruction converter to convert the instruction into one or more instructions of a different instruction set executable by the plurality of cores.

**30.** The apparatus of claim 21, wherein the operations include to check whether the first matrix is one of a pair of matrices, wherein N is 8 and M is 16, and wherein the plurality of cores comprise graphics cores.

**31.** An apparatus comprising: circuitry to convert a first instruction into one or more other instructions, the first instruction to indicate a first matrix having M rows by K columns of 32-bit elements each having a plurality of 4-bit data elements, a second matrix having K rows by N columns of 32-bit elements each having a plurality of 4-bit data elements, and a third matrix having M rows by N columns, wherein the 4-bit data elements of the first and second matrices are floating-point data elements; and execution circuitry to perform operations corresponding to the one or more other instructions, including to: for each data element position at a row m of the M rows and at a column n of the N columns of the third matrix: generate a dot product of the plurality of 4-bit data elements of each of K 32-bit elements of a row m of the first matrix and corresponding ones of the plurality of 4-bit data elements of each of K 32-bit elements of a column n of the second

matrix; and store a 32-bit result data element at the data element position of the third matrix, the 32-bit result data element to be based on accumulation of the dot product with a 32-bit data element corresponding to the data element position.

**32.** The apparatus of claim 31, wherein the operations include to check whether the first matrix is one of a pair of matrices.

**33.** The apparatus of claim 32, wherein N is 8 and M is 16, wherein the first and second matrices are to be stored in a plurality of registers, and further comprising graphics cores comprising the execution circuitry.

**34.** An apparatus comprising: an instruction converter to convert a first instruction into one or more other instructions, the first instruction to indicate a first matrix having M rows by K columns of 32-bit elements each having a plurality of 4-bit data elements, a second matrix having K rows by N columns of 32-bit elements each having a plurality of 4-bit data elements, and a third matrix having M rows by N columns, wherein the 4-bit data elements of the first and second matrices are floating-point data elements; and execution circuitry to perform operations corresponding to the one or more other instructions, including to: for each data element position at a row m of the M rows and at a column n of the N columns of the third matrix: generate a dot product of the plurality of 4-bit data elements of each of K 32-bit elements of a row m of the first matrix and corresponding ones of the plurality of 4-bit data elements of each of K 32-bit elements of a column n of the second matrix; and store a 32-bit result data element at the data element position of the third matrix, the 32-bit result data element to be based on accumulation of the dot product with a 32-bit data element corresponding to the data element position.

**35.** The apparatus of claim 34, wherein the operations include to check whether the first matrix is one of a pair of matrices.

**36.** The apparatus of claim 35, wherein N is 8 and M is 16, wherein the first and second matrices are to be stored in a plurality of registers, and further comprising graphics cores comprising the execution circuitry.

**37.** The apparatus of claim 35, wherein the instruction converter comprises a machine-readable storage medium storing code that when executed by the apparatus causes the apparatus to said convert the first instruction into the one or more other instructions.

**38.** A non-transitory machine-readable storage medium storing instructions that, when executed by a machine, cause the machine to perform operations, including to: receive an instruction that is to indicate a first matrix having M rows by K columns of 32-bit elements each having a plurality of 4-bit data elements, a second matrix having K rows by N columns of 32-bit elements each having a plurality of 4-bit data elements, and a third matrix having M rows by N columns, wherein the 4-bit data elements of the first and second matrices are floating-point data elements; and perform operations corresponding to the instruction, including to: for each data element position at a row m of the M rows and at a column n of the N columns of the third matrix: generate a dot product of the plurality of 4-bit data elements of each of K 32-bit elements of a row m of the first matrix and corresponding ones of the plurality of 4-bit data elements of each of K 32-bit elements of a column n of the second matrix; and store a 32-bit result data element at the data element position of the third matrix, the 32-bit result data element to be based on accumulation of the dot product with a 32-bit data element corresponding to the data element position.

**39.** The non-transitory machine-readable storage medium of claim 38, wherein the operations further include to check whether the first matrix is one of a pair of matrices.

---