

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265101

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Desai; Ranjit et al.

Method and Device for Process Data Sharing

Abstract

In one implementation, a method of accessing shared data among processes is performed by a device including processor(s), non-transitory memory, and an image acquisition interface. The method includes obtaining image data acquired by the image acquisition interface. The method further includes determining pose data based at least in part on inertial measurement unit (IMU) information measured by the image acquisition interface. The method also includes determining a gaze estimation based at least in part on eye tracking information obtained through the image acquisition interface. Based at least in part on characteristics of processes, the method includes determining an arrangement for the image data, the pose data, and the gaze estimation. The method additionally includes determining an access schedule for the processes based at least in part on at least one of: the arrangement, the characteristics of the processes, and hardware timing parameters associated with the device.

Inventors: Desai; Ranjit (Cupertino, CA), Rockwell; Michael J. (Palo Alto, CA), Duggineni; Venu Madhav (San Jose, CA), Lee; Robert Seon Wai (San Jose, CA)

Applicant: Apple Inc. (Cupertino, CA)

Family ID: 1000008578193

Appl. No.: 19/198802

Filed: May 05, 2025

Related U.S. Application Data

parent US continuation 17527923 20211116 parent-grant-document US 12307269 child US 19198802

parent US continuation 16548065 20190822 parent-grant-document US 11204783 child US 17527923

us-provisional-application US 62721685 20180823

Publication Classification

Int. Cl.: **G06F9/451** (20180101); **G06F3/01** (20060101); **G06F3/038** (20130101); **G06F9/48** (20060101); **G06F9/50** (20060101); **G06V40/16** (20220101); **G06V40/20** (20220101)

U.S. Cl.:

CPC **G06F9/451** (20180201); **G06F3/011** (20130101); **G06F3/013** (20130101); **G06F3/038** (20130101); **G06F9/4881** (20130101); **G06F9/5038** (20130101); **G06V40/166** (20220101); **G06V40/23** (20220101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 17/527,923, filed on Nov. 16, 2021, which is a continuation of U.S. patent application Ser. No. 16/548,065, filed on Aug. 22, 2019, which claims the benefit of U.S. Provisional App. No. 62/721,685, filed on Aug. 23, 2018, which are all hereby incorporated by reference herein in their entirety.

TECHNICAL FIELD

[0002] The present disclosure generally relates to data sharing, and in particular, to systems, methods, and devices providing low-latency data sharing using a deterministic pull/fetch model.

BACKGROUND

[0003] A physical environment refers to a physical world that people can sense and/or interact with without aid of electronic systems. Physical environments, such as a physical park, include physical articles, such as physical trees, physical buildings, and physical people. People can directly sense and/or interact with the physical environment, such as through sight, touch, hearing, taste, and smell.

[0004] In contrast, a computer-generated reality (CGR) environment refers to a wholly or partially simulated environment that people sense and/or interact with via an electronic system. In CGR, a subset of a person's physical motions, or representations thereof, are tracked, and, in response, one or more characteristics of one or more virtual objects simulated in the CGR environment are adjusted in a manner that comports with at least one law of physics. For example, a CGR system may detect a person's head turning and, in response, adjust graphical content and an acoustic field presented to the person in a manner similar to how such views and sounds would change in a physical environment. In some situations (e.g., for accessibility reasons), adjustments to characteristic(s) of virtual object(s) in a CGR environment may be made in response to representations of physical motions (e.g., vocal commands).

[0005] A person may sense and/or interact with a CGR object using any one of their senses, including sight, sound, touch, taste, and smell. For example, a person may sense and/or interact with audio objects that create 3D or spatial audio environment that provides the perception of point audio sources in 3D space. In another example, audio objects may enable audio transparency, which selectively incorporates ambient sounds from the physical environment with or without computer-generated audio. In some CGR environments, a person may sense and/or interact only with audio objects.

[0006] Examples of CGR include virtual reality and mixed reality.

[0007] A virtual reality (VR) environment refers to a simulated environment that is designed to be based entirely on computer-generated sensory inputs for one or more senses. A VR environment comprises a plurality of virtual objects with which a person may sense and/or interact. For

example, computer-generated imagery of trees, buildings, and avatars representing people are examples of virtual objects. A person may sense and/or interact with virtual objects in the VR environment through a simulation of the person's presence within the computer-generated environment, and/or through a simulation of a subset of the person's physical movements within the computer-generated environment.

[0008] In contrast to a VR environment, which is designed to be based entirely on computer-generated sensory inputs, a mixed reality (MR) environment refers to a simulated environment that is designed to incorporate sensory inputs from the physical environment, or a representation thereof, in addition to including computer-generated sensory inputs (e.g., virtual objects). On a virtuality continuum, a mixed reality environment is anywhere between, but not including, a wholly physical environment at one end and virtual reality environment at the other end.

[0009] In some MR environments, computer-generated sensory inputs may respond to changes in sensory inputs from the physical environment. Also, some electronic systems for presenting an MR environment may track location and/or orientation with respect to the physical environment to enable virtual objects to interact with real objects (that is, physical articles from the physical environment or representations thereof). For example, a system may account for movements so that a virtual tree appears stationery with respect to the physical ground.

[0010] Examples of mixed realities include augmented reality and augmented virtuality.

[0011] An augmented reality (AR) environment refers to a simulated environment in which one or more virtual objects are superimposed over a physical environment, or a representation thereof. For example, an electronic system for presenting an AR environment may have a transparent or translucent display through which a person may directly view the physical environment. The system may be configured to present virtual objects on the transparent or translucent display, so that a person, using the system, perceives the virtual objects superimposed over the physical environment. Alternatively, a system may have an opaque display and one or more imaging sensors that capture images or video of the physical environment, which are representations of the physical environment. The system composites the images or video with virtual objects, and presents the composition on the opaque display. A person, using the system, indirectly views the physical environment by way of the images or video of the physical environment, and perceives the virtual objects superimposed over the physical environment. As used herein, a video of the physical environment shown on an opaque display is called "pass-through video," meaning a system uses one or more image sensor(s) to capture images of the physical environment, and uses those images in presenting the AR environment on the opaque display. Further alternatively, a system may have a projection system that projects virtual objects into the physical environment, for example, as a hologram or on a physical surface, so that a person, using the system, perceives the virtual objects superimposed over the physical environment.

[0012] An augmented reality environment also refers to a simulated environment in which a representation of a physical environment is transformed by computer-generated sensory information. For example, in providing pass-through video, a system may transform one or more sensor images to impose a select perspective (e.g., viewpoint) different than the perspective captured by the imaging sensors. As another example, a representation of a physical environment may be transformed by graphically modifying (e.g., enlarging) portions thereof, such that the modified portion may be representative but not photorealistic versions of the originally captured images. As a further example, a representation of a physical environment may be transformed by graphically eliminating or obfuscating portions thereof.

[0013] An augmented virtuality (AV) environment refers to a simulated environment in which a virtual or computer generated environment incorporates one or more sensory inputs from the physical environment. The sensory inputs may be representations of one or more characteristics of the physical environment. For example, an AV park may have virtual trees and virtual buildings, but people with faces photorealistically reproduced from images taken of physical people. As

another example, a virtual object may adopt a shape or color of a physical article imaged by one or more imaging sensors. As a further example, a virtual object may adopt shadows consistent with the position of the sun in the physical environment.

[0014] There are many different types of electronic systems that enable a person to sense and/or interact with various CGR environments. Examples include smartphones, tablets, desktop/laptop computers, head-mounted systems, projection-based systems, heads-up displays (HUDs), vehicle windshields having integrated display capability, windows having integrated display capability, displays formed as lenses designed to be placed on a person's eyes (e.g., similar to contact lenses), headphones/earphones, speaker arrays, input systems (e.g., wearable or handheld controllers with or without haptic feedback and/or cameras having hand tracking and/or other body pose estimation abilities).

[0015] A head-mounted system may have one or more speaker(s) and an integrated opaque display. Alternatively, a head-mounted system may be a head-mounted enclosure (HME) configured to accept an external opaque display (e.g., a smartphone). The head-mounted system may incorporate one or more imaging sensors to capture images or video of the physical environment, and/or one or more microphones to capture audio of the physical environment. Rather than an opaque display, a head-mounted system may have a transparent or translucent display. The transparent or translucent display may have a medium through which light representative of images is directed to a person's eyes. The display may utilize digital light projection, OLEDs, LEDs, uLEDs, liquid crystal on silicon, laser scanning light source, or any combination of these technologies. The medium may be an optical waveguide, a hologram medium, an optical combiner, an optical reflector, or any combination thereof. In one implementation, the transparent or translucent display may be configured to become opaque selectively. Projection-based systems may employ retinal projection technology that projects graphical images onto a person's retina. Projection systems also may be configured to project virtual objects into the physical environment, for example, as a hologram or on a physical surface.

[0016] Content available on CGR devices is becoming more immersive, more graphically intensive, and universally applicable to everyday lives. Thus, the hardware in CGR devices continues to evolve to accommodate resource-heavy processes in order to keep up with the CGR content. However, with multiple processes contending for resources at once, latency and a large number of interrupts naturally create a bottleneck effect. Visible latency issues can adversely affect a user's experience.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0017] So that the present disclosure can be understood by those of ordinary skill in the art, a more detailed description may be had by reference to aspects of some illustrative implementations, some of which are shown in the accompanying drawings.

[0018] FIG. 1 is a block diagram of an exemplary operating environment in accordance with some implementations.

[0019] FIG. 2 is a block diagram illustrating an interrupt model for CGR data sharing in accordance with some implementations.

[0020] FIG. 3 is a block diagram of an exemplary deterministic pull/fetch model for CGR data sharing among processes in accordance with some implementations.

[0021] FIG. 4 is a block diagram illustrating an exemplary CGR data sharing process based on a deterministic pull/fetch model in accordance with some implementations.

[0022] FIG. 5A is a block diagram of an example of a controller that is used in a pull/fetch model in accordance with some implementations.

[0023] FIG. 5B is a block diagram illustrating an exemplary controller managed data arrangement and process access scheduling in accordance with some implementations.

[0024] FIG. 6 is a block diagram of an exemplary CGR device in accordance with some implementations.

[0025] FIG. 7A illustrates a CGR pipeline in accordance with some implementations.

[0026] FIG. 7B illustrates a distributed system implementing the CGR pipeline in accordance with some implementations.

[0027] FIG. 8 is a flowchart representation of a method of accessing shared data among processes in accordance with some implementations.

[0028] In accordance with common practice the various features illustrated in the drawings may not be drawn to scale. Accordingly, the dimensions of the various features may be arbitrarily expanded or reduced for clarity. In addition, some of the drawings may not depict all of the components of a given system, method or device. Finally, like reference numerals may be used to denote like features throughout the specification and figures.

SUMMARY

[0029] Various implementations disclosed herein include devices, systems, and methods for accessing shared data among processes. In various implementations, the method is performed at a device including one or more processors, non-transitory memory, and an image acquisition interface. The method includes obtaining image data associated with a field of view acquired by the image acquisition interface. The method further includes determining pose data based at least in part on inertial measurement unit (IMU) information, where the pose data corresponds to a current posture of the user measured by the image acquisition interface. The method additionally includes determining a gaze estimation based at least in part on eye tracking information obtained through the image acquisition interface. The method further also includes determining an arrangement for the image data, the pose data, and the gaze estimation based at least in part on a plurality of characteristics of a plurality of processes communicable with the image acquisition interface. The method further includes determining an access schedule for the plurality of processes based at least in part on at least one of: the arrangement for the image data, the pose data, and the gaze estimation, the plurality of characteristics of the plurality of processes, and hardware timing parameters associated with the device.

[0030] In accordance with some implementations, a device includes one or more processors, a non-transitory memory, and one or more programs; the one or more programs are stored in the non-transitory memory and configured to be executed by the one or more processors and the one or more programs include instructions for performing or causing performance of any of the methods described herein. In accordance with some implementations, a non-transitory computer readable storage medium has stored therein instructions, which, when executed by one or more processors of a device, cause the device to perform or cause performance of any of the methods described herein. In accordance with some implementations, a device includes: one or more processors, a non-transitory memory, and means for performing or causing performance of any of the methods described herein.

DESCRIPTION

[0031] Numerous details are described in order to provide a thorough understanding of the example implementations shown in the drawings. However, the drawings merely show some example aspects of the present disclosure and are therefore not to be considered limiting. Those of ordinary skill in the art will appreciate that other effective aspects and/or variants do not include all of the specific details described herein. Moreover, well-known systems, methods, components, devices and circuits have not been described in exhaustive detail so as not to obscure more pertinent aspects of the example implementations described herein.

[0032] As described above, processes on CGR devices following an interrupt model for resources often contend for resources simultaneously. As a result, the bottleneck effect may cause latency that

adversely affect the user's experience. Various implementations disclosed herein move away from the interrupt model towards a deterministic pull/fetch model. An access schedule is determined that allows the processes to fetch data at set times. The access schedule is determined based on known (e.g., deterministic) information, such as system parameters and user pose information. Accordingly, various implementations described herein address the above mentioned shortfalls, specifically those involved in data access. As a result, fewer memory and processing resources are consumed. It naturally follows that because latencies are decreased when data are through the CGR display pipeline, the overall user experience can be improved.

[0033] FIG. 1 is a block diagram of an exemplary operating environment **100** in accordance with some implementations. While pertinent features are shown, those of ordinary skill in the art will appreciate from the present disclosure that various other features have not been illustrated for the sake of brevity and so as not to obscure more pertinent aspects of the example implementations disclosed herein. To that end, as a non-limiting example, the operating environment **100** includes a controller **102** and a CGR device **104**.

[0034] In some implementations, the CGR device **104** corresponds to tablet or mobile phone. In various implementations, the CGR device **104** corresponds to a head-mounted system, such as a head-mounted device (HMD) or a head-mounted enclosure (HME) having a tablet or mobile phone inserted therein. In some implementations, the CGR device **104** is configured to present CGR content to a user. In some implementations, the CGR device **104** includes a suitable combination of software, firmware, and/or hardware.

[0035] According to some implementations, the CGR device **104** presents, via a display **122**, CGR content to the user while the user is virtually and/or physically present within a scene **106** that includes a table **107** within the field-of-view **111** of the CGR device **104**. In some implementations, the CGR device **104** is configured to present virtual content (e.g., the virtual cylinder **109**) and to enable video pass-through of the scene **106** (e.g., including a representation **117** of the table **107**) on a display **122**. In some implementations, the CGR device **104** is configured to present virtual content and to enable optical see-through of the scene **106**.

[0036] In some implementations, the user holds the CGR device **104** in his/her hand(s). In some implementations, the user wears the CGR device **104** on his/her head. As such, the CGR device **104** includes one or more CGR displays provided to display the CGR content. For example, the CGR device **104** encloses the field-of-view of the user. In some implementations, the CGR device **104** is replaced with a CGR chamber, enclosure, or room configured to present CGR content in which the user does not wear the CGR device **104**.

[0037] In some implementations, the controller **102** is configured to manage and coordinate presentation of CGR content for the user. In some implementations, the controller **102** includes a suitable combination of software, firmware, and/or hardware. In some implementations, the controller **102** is a computing device that is local or remote relative to the scene **106**. For example, the controller **102** is a local server located within the scene **106**. In another example, the controller **102** is a remote server located outside of the scene **106** (e.g., a cloud server, central server, etc.). In some implementations, the controller **102** is communicatively coupled with the CGR device **104** via one or more wired or wireless communication channels **144** (e.g., BLUETOOTH, IEEE 802.11x, IEEE 802.16x, IEEE 802.3x, etc.). In some implementations, the functionalities of the controller **102** are provided by and/or combined with the CGR device **104**.

[0038] As illustrated in FIG. 1, the CGR device **104** presents a representation of the scene **106**. In some implementations, the representation of the scene **106** is generated by the controller **102** and/or the CGR device **104**. In some implementations, the representation of the scene **106** includes a virtual scene that is a simulated replacement of the scene **106**. In other words, in some implementations, the representation of the scene **106** is simulated by the controller **102** and/or the CGR device **104**. In such implementations, the representation of the scene **106** is different from the scene **106** where the CGR device **104** is located. In some implementations, the representation of the

scene **106** includes an augmented scene that is a modified version of the scene **106** (e.g., including the virtual cylinder **109**). For example, in some implementations, the controller **102** and/or the CGR device **104** modify (e.g., augment) an image of the scene **106** in order to generate the representation of the scene **106**. In some implementations, the controller **102** and/or the CGR device **104** generate the representation of the scene **106** by simulating a replica of the scene **106**. In some implementations, the controller **102** and/or the CGR device **104** generate the representation of the scene **106** by removing and/or adding items from the simulated replica of the scene **106**.

[0039] FIG. 2 is a block diagram illustrating an interrupt model **200** for CGR data sharing. In some implementations, in order to generate the representation of the scene **106** (as illustrated in FIG. 1), a plurality of processes, e.g., process 1 **210-1**, process 2 **210-2**, process 3 **210-3** . . . process N **210-N**, obtain raw data acquired by the CGR device **104** (as illustrated in FIG. 1) as inputs. For example, the raw data includes image data **230** acquired by image sensor(s), pose data **240** acquired by an IMU, gaze estimation data **250** derived from information obtained by an eye tracker, and other data **260** acquired by the CGR device **104**. In some implementations, the plurality of processes **210** are dependent upon each other, such that outputs from one process are used by another process as inputs, e.g., outputs from process 2 **210-2** are inputs to process 3 **210-3**. In some implementation, the raw data and/or the outputs from processes **210** are communicated through a communication path **220**, e.g., a communication path established through communication interface(s) of the controller **102** and/or the communication interface(s) of the CGR device **104**.

[0040] As shown in FIG. 2, when process 1 **210-1** needs the image data **230** and the gaze estimation data **250**, process 1 **210-1** interrupts the image sensor and the eye tracker in order to obtain a copy of the image data **230** and a copy of the gaze estimation data **250** as inputs. Likewise, when process 2 **210-2** also needs the image data **230** and the gaze estimation data **250**, process 2 **210-2** interrupts the image sensor and the eye tracker in order to obtain a copy of the image data **230** and a copy of the pose estimation data **250** as inputs. In another example, as shown in FIG. 2, process 3 **210-3** needs inputs from process 2 **210-2**, the pose data **240**, and the other data **260**. Process 3 **210-3** would wait for the completion of process 2 **210-2** while interrupting the IMU and other sensor(s) in order to obtain the pose data **240** and the other data **260**.

[0041] The interrupt model **200** is inefficient for several reasons. First, when multiple processes (e.g., process 1 **210-1** and process 2 **210-2**) are contending for resources (e.g., the image data **230** and the gaze estimation data **250**), the interrupts created multiple bottlenecks, e.g., at least one bottleneck at the sensors and another at the communication path **220**. Second, because multiple copies of the data are created for multiple processes, the interrupt model **200** does not share memory across different tasks. As such, the memory usage is inefficient. Due to the bottlenecks, the inefficient memory utilization, and the cascade effect from process dependencies, the interrupt model **200** cannot meet the latency requirement for real-time streaming of CGR content. As such, the CGR scene presented using the interrupt model **200** can cause motion sickness for a user.

[0042] FIG. 3 is a block diagram of an example deterministic pull/fetch model **300** for CGR data sharing among processes in accordance with some embodiments. A deterministic system typically involves no randomness in the development of future states of the system. A deterministic model will thus produce the same output from a given starting condition or an initial state. As such, using the deterministic model **300**, a system can predict when data from where would be produced and accessed. As such, the pull/fetch model **300** allows processes to retrieve data when the data are ready to be fetched and from a location efficient for the retrieval. Accordingly, in such system, bottlenecks are reduced and contention for resources goes down.

[0043] For example, FIG. 3 shows two processes, namely process M and process N, where process N depends on outputs from process M and process M fetches image data taken by a camera. The system learns hardware timing parameters, such as camera exposure time $\Delta 1$, as well as characteristics of the processes, e.g., processing time $\Delta 2$ for process M and/or the type of input data for processes M and N. These parameters and characteristics are used by the system to determine

when a process would fetch data and from which location the data would be fetched. In FIG. 3, the system uses a system synchronization clock to measure the starting time (e.g., T.sub.1), the expected camera operation duration (e.g., from T.sub.2 to T.sub.3), and the expected process M execution time (from T.sub.4 to T.sub.5) and calculates the wakeup time T.sub.4 for process M and wakeup time T.sub.6 for process N.

[0044] In some embodiments, the deterministic system generates an access schedule for the processes including the calculated wakeup time for each process. In some embodiments, between each task, a threshold amount of time is reserved, e.g., the periods between T.sub.1 and T.sub.2, between T.sub.3 and T.sub.4, and between T.sub.5 and T.sub.6. The threshold amount of time is reserved for communication latency between processes and/or hardware components, e.g., the time between T.sub.1 and T.sub.2 is for the system notifying the camera to start the image data acquisition, the time between T.sub.3 and T.sub.4 is for the camera waking up process M, and the time between T.sub.5 and T.sub.6 is for process M waking up process N. In some embodiments, the threshold amount of time is also recorded in the access schedule. Following the access schedule, upon waking up, the inputs for processes M and N are ready to be fetched. Thus, in contrast to the interrupt model **200** as shown in FIG. 2, the pull/fetch model **300** as shown in FIG. 3 reduces constraints at a system level and increases system performance as a whole.

[0045] FIG. 4 is a block diagram illustrating a CGR data sharing process **400** based on the deterministic pull/fetch model **300** in accordance with some implementations. In some implementations, in order to generate the representation of the scene **106** (as illustrated in FIG. 1), a plurality of processes, e.g., process **1 410-1**, process **2 410-2**, process **3 410-3** . . . process **N 410-N**, receive raw data acquired by the CGR device **104** (as illustrated in FIG. 1). For example, the raw data include image data **430** acquired by image sensor(s), pose data **440** acquired by an IMU, gaze estimation data **450** derived from information obtained by an eye tracker, and other data **460** acquired by the CGR device **104**. In some implementations, the plurality of processes are dependent upon each other, such that outputs from one process are used by another process as inputs, e.g., outputs from process **2 410-2** are inputs to process **3 410-3**. In some implementation, the raw data and/or the outputs from the plurality of processes **410** are communicated through a communication path **420**, e.g., a communication path established through communication interface(s) of the controller **102** and/or the communication interface(s) of the CGR device **104**.

[0046] As shown in FIG. 4, using the deterministic model **300**, the controller **102** directs data arrangements (e.g., layout and/or sequencing) for data **427** in a buffer **425** and wakes up the plurality of processes **410** according to an access schedule, e.g., when the data in the buffer **425** are ready to be pulled/fetched. For example, the image data **427-2** for process **2 410-2** (represented as (I, 2) in FIG. 4) and the image data **427-4** for process **3 410-3** (represented as (I, 3) in FIG. 4) are obtained from the camera captured image data **430**. The image data **427-2** for process **2 410-2** and the image data **427-4** for process **3 410-3** are arranged in the buffer **425** in a layout and/or sequence such that when the image data **427-2** and **427-4** are provided to the processes **410** in a data stream, and process **2 410-2** and process **3 410-3** wake up in sequence according to the access schedule. In particular, when process **3 410-3** wakes up, the image data **427-4** as well as the outputs from process **2 410-2** are ready to be pulled/fetched.

[0047] FIG. 5A is a block diagram of an example of the controller **102** that is used in the pull/fetch model described above in accordance with some implementations. FIG. 5B is a block diagram illustrating the data arrangement and process access scheduling managed by the controller **102** in accordance with some implementations. FIG. 5B is used to illustrate an exemplary CGR scene generation managed by the controller **102** shown in FIG. 5A in accordance with some implementations. While certain specific features are illustrated, those skilled in the art will appreciate from the present disclosure that various other features have not been illustrated for the sake of brevity, and so as not to obscure more pertinent aspects of the implementations disclosed herein. To that end, as a non-limiting example, as shown in FIG. 5A, in some implementations the

controller **102** includes one or more processing units **502** (e.g., microprocessors, application-specific integrated-circuits (ASICs), field-programmable gate arrays (FPGAs), graphics processing units (GPUs), central processing units (CPUs), processing cores, and/or the like), one or more input/output (I/O) devices **506**, one or more communication interfaces **508** (e.g., universal serial bus (USB), FIREWIRE, THUNDERBOLT, IEEE 802.3x, IEEE 802.11x, IEEE 802.16x, global system for mobile communications (GSM), code division multiple access (CDMA), time division multiple access (TDMA), global positioning system (GPS), infrared (IR), BLUETOOTH, ZIGBEE, and/or the like type interface), one or more programming (e.g., I/O) interfaces **510**, a memory **520**, and one or more communication buses **504** for interconnecting these and various other components. [0048] In some implementations, the one or more communication buses **504** include circuitry that interconnects and controls communications between system components. In some implementations, the one or more I/O devices **506** include at least one of a keyboard, a mouse, a touchpad, a joystick, one or more microphones, one or more speakers, one or more image sensors, one or more displays, and/or the like.

[0049] The memory **520** includes high-speed random-access memory, such as dynamic random-access memory (DRAM), static random-access memory (SRAM), double-data-rate random-access memory (DDR RAM), or other random-access solid-state memory devices. In some implementations, the memory **520** includes non-volatile memory, such as one or more magnetic disk storage devices, optical disk storage devices, flash memory devices, or other non-volatile solid-state storage devices. The memory **520** optionally includes one or more storage devices remotely located from the one or more processing units **502**. The memory **520** comprises a non-transitory computer readable storage medium. In some implementations, the memory **520** or the non-transitory computer readable storage medium of the memory **520** stores the following programs, modules and data structures, or a subset thereof including an optional operating system **530** and a CGR content module **540**.

[0050] The operating system **530** includes procedures for handling various basic system services and for performing hardware dependent tasks. In some implementations, the CGR content module **540** is configured to manage and coordinate presentation of CGR content for one or more users (e.g., a single set of CGR content for one or more users, or multiple sets of CGR content for respective groups of one or more users). To that end, in various implementations, the CGR content module **540** includes a data obtaining unit **542**, a tracking unit **544**, a coordination unit **546**, a data transmitting unit **548**, and a data access unit **550**.

[0051] In some implementations, the data obtaining unit **542** is configured to obtain data (e.g., image data, pose data, gaze estimation, presentation data, interaction data, sensor data, location data, etc.) from at least the CGR device **104**. To that end, in various implementations, the data obtaining unit **542** includes instructions and/or logic therefor, and heuristics and metadata therefor.

[0052] In some implementations, the tracking unit **544** is configured to map the scene **106** and to track the position/location of at least the CGR device **104** with respect to the scene **106**. To that end, in various implementations, the tracking unit **544** includes instructions and/or logic therefor, and heuristics and metadata therefor.

[0053] In some implementations, the coordination unit **546** is configured to manage and coordinate the presentation of CGR content to the user by the CGR device **104**. In order to manage and coordinate the presentation of CGR content, in some implementations, the coordination unit **546** is configured to obtain information related to processes, data storage, and hardware characteristics. The information is then used by the coordination unit **546** to coordinate the processes at the system level. To that end, in various implementations, the coordination unit **546** includes instructions and/or logic therefor, and heuristics and metadata therefor.

[0054] In some implementations, the data transmitting unit **548** is configured to transmit data (e.g., presentation data, location data, etc.) to at least the CGR device **104**. To that end, in various implementations, the data transmitting unit **548** includes instructions and/or logic therefor, and

heuristics and metadata therefor.

[0055] In some implementations, the data access unit **550** is configured to determine an access schedule that allows processes to fetch data obtained from the data obtaining unit **542**. In some implementations, the data access unit **550** determines a data arrangement **552** for efficient data retrieval by the processes. In some implementations, the data access unit **550** also determines an access schedule **554** for the processes to pull or fetch the data. To that end, in various implementations, the data transmitting unit **548** includes instructions and/or logic therefor, and heuristics and metadata therefor.

[0056] For example, in FIG. 5B, a scene recognition process **560** obtains image data and pose data **580** for recognizing the scene **106**. As an eye tracker detects a user fixing gaze proximate to a region of interest (ROI) **572** within the scene **106**, a ROI recognition process **562** identifies ROI image data **590** in the image data based on gaze estimation data **582**. Subsequently, an object recognition process **564** analyzes the ROI image data **590** and recognizes an object **574** within the ROI **572**. As shown in FIG. 5B, the processes **560**, **562**, and **564** access the data arrangement **552** according to the access schedule **554**, where the access schedule **554** includes information such as a starting time for each process. Also as shown in FIG. 5B, data in the data arrangement **552** are arranged to accommodate the scheduled fetching. For instance, since the ROI image data **590** are used by the ROI recognition process **562** and the object recognition process **564**, the ROI image data **590** is stored separately from non-ROI image data **592** to accommodate more frequent or more urgent access.

[0057] Referring back to FIG. 5A, although the data obtaining unit **542**, the tracking unit **544**, the coordination unit **546**, the data transmitting unit **548**, and the data access unit **550** are shown as residing on a single device (e.g., the controller **102**), it should be understood that in other implementations, any combination of the data obtaining unit **542**, the tracking unit **544**, the coordination unit **546**, the data transmitting unit **548**, and the data access scheduling unit **550** may be located in separate computing devices.

[0058] Moreover, FIG. 5A is intended more as functional description of the various features which are present in a particular embodiment as opposed to a structural schematic of the implementations described herein. As recognized by those of ordinary skill in the art, items shown separately could be combined and some items could be separated. For example, some functional modules shown separately in FIG. 5A could be implemented in a single module and the various functions of single functional blocks could be implemented by one or more functional blocks in various implementations. The actual number of modules and the division of particular functions and how features are allocated among them will vary from one embodiment to another and, in some implementations, depends in part on the particular combination of hardware, software, and/or firmware chosen for a particular embodiment.

[0059] FIG. 6 is a block diagram of an exemplary CGR device **104** in accordance with some implementations. While certain specific features are illustrated, those skilled in the art will appreciate from the present disclosure that various other features have not been illustrated for the sake of brevity, and so as not to obscure more pertinent aspects of the implementations disclosed herein. To that end, as a non-limiting example, in some implementations the CGR device **104** includes one or more processing units **602** (e.g., microprocessors, ASICs, FPGAs, GPUs, CPUs, processing cores, and/or the like), one or more input/output (I/O) devices **606**, one or more communication interfaces **608** (e.g., USB, FIREWIRE, THUNDERBOLT, IEEE 802.3x, IEEE 802.11x, IEEE 802.16x, GSM, CDMA, TDMA, GPS, IR, BLUETOOTH, ZIGBEE, and/or the like type interface), one or more programming (e.g., I/O) interfaces **610**, one or more CGR displays **612**, one or more image acquisition interfaces **614** (e.g., optional interior and/or exterior facing image sensors), a memory **620**, and one or more communication buses **604** for interconnecting these and various other components.

[0060] In some implementations, the one or more communication buses **604** include circuitry that

interconnects and communications between system components. In some implementations, the one or more I/O devices **606** include at least one of one or more physiological sensors (e.g., blood pressure monitor, heart rate monitor, blood oxygen sensor, blood glucose sensor, etc.), one or more microphones, one or more speakers, a haptics engine, and/or the like.

[0061] In some implementations, the one or more CGR displays **612** are configured to present CGR content to the user. In some embodiments, the one or more CGR displays **612** correspond to holographic, digital light processing (DLP), liquid-crystal display (LCD), liquid-crystal on silicon (LCoS), organic light-emitting field-effect transitory (OLET), organic light-emitting diode (OLED), surface-conduction electron-emitter display (SED), field-emission display (FED), quantum-dot light-emitting diode (QD-LED), micro-electro-mechanical system (MEMS), and/or the like display types. In some embodiments, the one or more CGR displays **612** correspond to diffractive, reflective, polarized, holographic, etc. waveguide displays. For example, the CGR device **104** includes a single AR/VR display. In another example, the CGR device **104** includes an CGR display for each eye of the user.

[0062] In some implementations, the one or more image acquisition interfaces **614** are configured to obtain data for CGR content generation. In some implementations, the one or more image acquisition interfaces **614** include at least one of one or more image sensors, an inertial measurement unit (IMU), an accelerometer, a gyroscope, a thermometer, one or more depth sensors (e.g., a structured light, a time-of-flight, or the like), an eye tracker. For example, the one or more image sensors correspond to one or more RGB camera (e.g., with a complimentary metal-oxide-semiconductor (CMOS) image sensor or a charge-coupled device (CCD) image sensor), IR camera, event-based camera, and/or the like.

[0063] The memory **620** includes high-speed random-access memory, such as DRAM, SRAM, DDR RAM, or other random-access solid-state memory devices. In some implementations, the memory **620** includes non-volatile memory, such as one or more magnetic disk storage devices, optical disk storage devices, flash memory devices, or other non-volatile solid-state storage devices. The memory **620** optionally includes one or more storage devices remotely located from the one or more processing units **602**. The memory **620** comprises a non-transitory computer readable storage medium. In some implementations, the memory **620** or the non-transitory computer readable storage medium of the memory **620** stores the following programs, modules and data structures, or a subset thereof including an optional operating system **630** and a CGR presentation module **640**.

[0064] The operating system **630** includes procedures for handling various basic system services and for performing hardware dependent tasks. In some implementations, the CGR presentation module **640** is configured to present CGR content to the user via the one or more CGR displays **612**. To that end, in various implementations, the CGR presentation module **640** includes a data obtaining unit **642**, a CGR presenting unit **644**, an eye tracking unit **646**, and a data transmitting unit **648**.

[0065] In some implementations, the data obtaining unit **642** is configured to obtain data (e.g., image data, pose data, presentation data, interaction data, sensor data, location data, etc.) from at least the controller **102**. To that end, in various implementations, the data obtaining unit **642** includes instructions and/or logic therefor, and heuristics and metadata therefor.

[0066] In some implementations, the CGR presenting unit **644** is configured to present CGR content via the one or more CGR displays **612**. To that end, in various implementations, the CGR presenting unit **644** includes instructions and/or logic therefor, and heuristics and metadata therefor. In some implementations, the CGR presenting unit **644** is configured to project an image comprising emitted light in a first wavelength range through an eyepiece that distorts light in the first wavelength range. In some embodiments, the CGR presenting unit **644** is configured to project an image comprising emitted light in a first wavelength through an eyepiece that reflects and refracts light in the first wavelength range while passing, without substantial distortion, light in the

second wavelength range.

[0067] In some implementations, the eye tracking unit **646** is configured to emit, using one or more light sources disposed between the eyepiece and the display, light in a second wavelength range and detect, using a camera, the light in the second wavelength range. In various implementations, the one or more light sources illuminate the eye of a user and the camera detect light reflected from the eye of the user. To that end, in various implementations, the eye tracking unit **646** includes instructions and/or logic therefor, and heuristics and metadata therefor. In some implementations, the eye tracking unit **646** is configured to emitting light in a second wavelength range through the eyepiece and detecting the light in the second wavelength range reflected by the eye of a user. In some implementations, the eye tracking unit **646** provides a gaze estimation based at least in part on the detected light reflection from the eye of the user.

[0068] In some implementations, the data transmitting unit **648** is configured to transmit data (e.g., presentation data, location data, etc.) to at least the controller **102**. To that end, in various implementations, the data transmitting unit **648** includes instructions and/or logic therefor, and heuristics and metadata therefor.

[0069] Although the data obtaining unit **642**, the CGR presenting unit **644**, the eye tracking unit **646**, and the data transmitting unit **648** are shown as residing on a single device (e.g., the CGR device **104**), it should be understood that in other implementations, any combination of the data obtaining unit **642**, the CGR presenting unit **644**, the eye tracking unit **646**, and the data transmitting unit **648** may be located in separate computing devices.

[0070] Moreover, FIG. **6** is intended more as functional description of the various features which are present in a particular embodiment as opposed to a structural schematic of the implementations described herein. As recognized by those of ordinary skill in the art, items shown separately could be combined and some items could be separated. In some implementations, some functional modules shown separately in FIG. **6** could be implemented in a single module and the various functions of single functional blocks could be implemented by one or more functional blocks in various implementations. The actual number of modules and the division of particular functions and how features are allocated among them will vary from one embodiment to another and, in some implementations, depends in part on the particular combination of hardware, software, and/or firmware chosen for a particular embodiment.

[0071] Additionally, in some implementations, the functions performed by the controller **102** as shown in FIG. **5A** and the CGR device **104** as shown in FIG. **6** are distributed among devices. For example, FIG. **7A** illustrates a CGR pipeline **700** that receives data from sensors and generates the representation of the scene **106** (as illustrated in FIG. **1**) in accordance with some implementations. FIG. **7B** illustrates a system **750** that implements the CGR pipeline **700**. The system **750** distributes tasks performed by the controller **102** and/or the CGR device **104** described above between the controller **102** and the CGR device **104**.

[0072] As shown in FIG. **7A**, in some implementations, the image acquisition interface **614** (FIG. **6**) includes at least an image sensor **710** for outputting image data, an IMU **712** for outputting pose data, an eye tracker **714** for providing gaze estimation, and one or more other sensors **716** for providing raw data as inputs to processes in order to generate the scene **106**. Upon receiving the image data, pose data, gaze estimation data, and raw data, processes executed by the controller **102** stores the received data in a buffer (e.g., the buffer **425** in FIG. **4** or the buffer **720** in FIG. **7A**) according to the data arrangement **552** as explained above. Further, the processes executed by the controller **102** accesses the buffer **720** according to the access schedule **554**. The access schedule **554** is generated according to the deterministic model as described above with reference to FIG. **3**.

[0073] In some implementations, the CGR pipeline **700** includes a rendering module **730** that receives CGR content and the data from the buffer **720** and renders an image on the display **612**. In various implementations, the CGR content includes definitions of geometric shapes of virtual objects, colors and/or textures of virtual objects, images (such as a see-through image of the scene

106), and other information describing content to be represented in the rendered image. In some implementations, final correction is performed prior to displaying the rendered image. For example, based on the pose data, the rendered image is corrected to improve the user's experience. [0074] In some implementations, the final correction and other less computationally-intensive tasks (e.g., sensor data preprocessing) are performed at the CGR device, as shown in FIG. 7B. The system **750** in FIG. 7B, which implements the CGR pipeline **700** according to some embodiments, distributes the functions performed by the controller **102** and/or the CGR device **104**. As such, more computationally-intensive tasks are performed at the controller **102**, e.g., using one or more processors **756** of the controller **102** for tasks such as machine learning, computer vision, and/or 3D rendering etc. The controller **102** then transports a computed image to the CGR device **104** for final correction before display. In some implementations, the transportation process includes compression/decompression and communications between the controller **102** and the CGR device **104**.

[0075] In some implementations, the one or more processors **756** includes the coordination unit **546** and the data access unit **550**, which further includes the data arrangement **552-1** and the access schedule **554-1** for processes distributed across the system **750**. For example, the controller **102** collects characteristics of processes and hardware parameters from the CGR device **104**. Using machine learning such as neural networks, characteristics of processes across platforms, including the transportation process, the sensor data preprocessing, and/or the final correction etc., can be extracted and weights are assigned, so that a sequencing of the processes accessing data can be predicted. The controller **102** can also determine the data arrangement **552-1** to accommodate the predicted data access based at least in part on the access schedule **554-1** in some implementations. Though FIG. 7B illustrates the system **750** comprising one controller **102** and one CGR device **104**, in some embodiments, multiple CGR devices can connect to the controller **102**. In such embodiments, the machine learning on the controller **102** can be used to predict access schedule for cross-platform processes including processes across multiple CGR devices (e.g., in a multi-player CGR environment). In some implementations, the controller **102** distributes the system wide data arrangement **552-1** and access schedule **554-1**, so that each CGR device obtains a copy or a subset of the data arrangement **552-1** and/or the access schedule **554-1** from the controller **102**. According to a local copy of the data arrangement **552-2** and a local copy of the access schedule **554-2**, data acquired by the image acquisition interface **614** are arranged and processes associated with the CGR device **104** access the arranged data according to the schedule. As such, different from other systems, where each process or each device may have been optimized for data access, the system **750** according to embodiments described herein provides platform agnostic scheduling for data access at system level, so that latency is reduced as a whole.

[0076] FIG. 8 is a flowchart representation of a method **800** of determining an access schedule for processes in accordance with some implementations. In various implementations, the method **800** is performed by a device with one or more processors, non-transitory memory, and an image acquisition interface (e.g., the image acquisition interface **614** in FIG. 6). In some implementations, the method **800** is performed by processing logic, including hardware, firmware, software, or a combination thereof. In some implementations, the method **800** is performed by a processor and/or a controller (e.g., the controller **102** in FIG. 1) executing instructions (e.g., code) stored in a non-transitory computer-readable medium (e.g., a memory). Briefly, in some circumstances, the method **800** includes: utilizing the image acquisition interface to obtain image data, determine pose data, and determine a gaze estimation; determining an arrangement for the image data, the pose data, and the gaze estimation based at least in part on a plurality of characteristics of a plurality of processes communicable with the image acquisition interface; and determining an access schedule for the plurality of processes based at least in part on at least one of: the arrangement for the image data, the pose data, and the gaze estimation, the plurality of characteristics of the plurality of processes, and hardware timing parameters associated with the device.

[0077] The method **800** begins, in block **810**, with the device obtaining image data associated with a field of view acquired by the image acquisition interface. In some embodiments, the image acquisition interface includes an image sensor (e.g., the image sensor **710** in FIG. 7A) for acquiring the image data associated with the field of view.

[0078] The method **800** continues, in block **820**, with the device determining pose data based at least in part on inertial measurement unit (IMU) information, where the pose data corresponds to a current posture of the user measured by the image acquisition interface. In some embodiments, the image acquisition interface includes an IMU (e.g., the IMU **712** in FIG. 7A) for measuring the current posture of the user.

[0079] The method **800** continues, in block **830**, with the device determining a gaze estimation based at least in part on eye tracking information obtained through the image acquisition interface. In some embodiments, the image acquisition interface includes an eye tracker (e.g., the eye tracker **714** in FIG. 7A) for obtaining the eye tracking information.

[0080] The method **800** continues, in block **840**, with the device determining an arrangement for the image data, the pose data, and the gaze estimation based at least in part on a plurality of characteristics of a plurality of processes communicable with the image acquisition interface. For example, as shown in FIG. 4, the arrangement of the data **427** in the buffer **425** is determined by the controller **102** using the deterministic model **300**. The deterministic model **300**, as shown in FIG. 3, obtains characteristics (e.g., execution sequencing and/or process execution duration) of processes as part of the parameters for determining the arrangement of the data **427** in the buffer **425**.

[0081] In some embodiments, as represented by block **842**, the arrangement determination includes aggregating data for different processes and arranging the aggregated data based at least in part on the plurality of characteristics of the plurality of processes. For instance, the arrangement determination for two processes includes the steps of: (1) obtaining at least one of a first image data, a first pose data, and a first gaze estimation for a first process of the plurality of processes, where the first process is characterized by a first set of characteristics; (2) obtaining at least one of a second image data, a second pose data, and a second gaze estimation for a second process of the plurality of processes, wherein the second process is characterized by a second set of characteristics; and (3) aggregating the first image data, the first pose data, and the first gaze estimation with the second image data, the second pose data, and the second gaze estimation to generate the image data, the pose data, and the gaze estimation, wherein the image data, the pose data, and the gaze estimation are arranged based at least in part on the first set of characteristics and the second set of characteristics.

[0082] For example, as shown in FIG. 4, the arrangement of data **427** in the buffer **425** is determined based at least in part on the characteristics of the processes **410**. In case the outputs from process **2 410-2** are used by process **3 410-3** as inputs, the data **427-2** for process **2 410-2** is arranged in the buffer such that it is ready to be fetched by process **2 410-2** before the data **427-4** is ready to be fetched by process **3 410-3**. In another example, as shown in FIG. 5B, the ROI image data **590** is used by more processes than the non-ROI image data **592**. As such, according to the data arrangement **552**, the ROI image data **590** is stored at a more frequently access region and/or a region for more urgent needs.

[0083] Still referring to FIG. 8, the method **800** continues, in block **850**, with the device determining an access schedule for the plurality of processes based at least in part on at least one of: the arrangement for the image data, the pose data, and the gaze estimation, the plurality of characteristics of the plurality of processes, and hardware timing parameters associated with the device. In some embodiments, as represented by block **852**, the method **800** includes determining the access schedule based on a deterministic model. In such embodiments, the method **800** further includes determining, for a first process of the plurality of processes, inputs for the first process, at least one of a hardware or a second process providing the inputs, and a time parameter for

obtaining the inputs from at least one of the hardware or the second process; and calculating a waking time for the first process based at least in part on the time parameter. For example, as shown in FIG. 3, for process M, the waking time for execution of process M is determined at least in part on the exposure time of the camera, e.g., known $\Delta 1$. For process N, the waking time for execution of process M is determined at least in part on the exposure time of the camera and the execution duration of process M, e.g., known $\Delta 1$ and known $\Delta 2$.

[0084] In some embodiments, as represented by block **854**, the method **800** includes distributing computation-intensive tasks to a base device and performing a final correction at the device in order to optimize the user experience (e.g., reduce motion sickness). For example, in the CGR pipeline **700** as shown in FIG. 7A, the computation-intensive tasks including rendering (as performed by the rendering module **730**) can be performed by the controller **102**, as shown in FIG. 7B; while minimal computation such as sensor data preprocessing can be performed by the CGR device **104**, as shown in FIG. 7B. Further as shown in FIG. 7A, prior to displaying the scene, data from the CGR device **104** (e.g., pose data) can be used for final correction. In some embodiments, the frames sent by the base for rendering are timestamped. Knowing the time when the frames are computed at the base, upon receiving the frames from the base, the CGR device predicts a rendering pose at a rendering time based on a trajectory of the pose data (e.g., corresponding to a current posture of the user as measured by the image acquisition interface). The CGR device then performs final correction by adjusting the frames using the rendering pose at the rendering time.

[0085] In some embodiments, as represented by block **856**, the method **800** includes determining the access schedule and/or the arrangement by a second device. In particular, the distributed access schedule determination includes the steps of triggering collection of data by a base device, wherein the base device obtains at least one of the arrangement for the image data, the pose data, and the gaze estimation, the plurality of characteristics of the plurality of processes, and the hardware timing parameters from the device, and the data are also collected by the base device from other devices; and receiving the access schedule from the second device, wherein the access schedule is determined by the second device based on the data. For example, as shown in FIG. 7B, the controller **102** obtains processes characteristics, data arrangement information, and hardware parameters etc. from the CGR device **104**. In some embodiments, the controller **102** also receives such information from other CGR devices. Utilizing the collected information, the controller **102** provides the system wide access schedule **554-1** and/or data arrangement **552-1**.

[0086] In some embodiments, the method **800** continues, in block **860**, with the device accessing the arrangement according to the access schedule in order to generate a scene for display (e.g., a CGR scene that is a representation of a real-world scene or a CGR scene that is a fully virtual scene); and displaying the scene using data obtained from the arrangement and according to the access schedule, where the data is a subset of at least one of the image data, the pose data, and the gaze estimation. For example, in FIG. 7A, the scene provided to the display **612** of the CGR device is generated using the data from the buffer **720**. As shown in FIG. 4, the processes **410** for producing the scene pulls the data stream from the buffer **425** according to the access schedule, where data **427** in the buffer **425** are arranged according to the arrangement. In another example, as shown in FIG. 5B, the scene recognition process **560**, the ROI recognition process **562**, and the object recognition process **564** access data according to the access schedule **554**, and the ROI image data **590** and the non-ROI image data **592** are stored according to the data arrangement **552**.

[0087] In some embodiments, as represented by block **862**, accessing the arrangement according to the access schedule includes determining an expected execution time for a process of the plurality of processes based on the access schedule and waking up the process for data access at the expected execution time upon fetching data used by the process according to the arrangement. For example, as shown in FIG. 3, based on the access schedule, the expected execution time for process M at time T_4 is determined according to the access scheduled. At time T_4 , the image data from the camera are fetched by the process M from a buffer, in which data are arranged according to the

system determined data arrangement.

[0088] While various aspects of implementations within the scope of the appended claims are described above, it should be apparent that the various features of implementations described above may be embodied in a wide variety of forms and that any specific structure and/or function described above is merely illustrative. Based on the present disclosure one skilled in the art should appreciate that an aspect described herein may be implemented independently of any other aspects and that two or more of these aspects may be combined in various ways. For example, an apparatus may be implemented and/or a method may be practiced using any number of the aspects set forth herein. In addition, such an apparatus may be implemented and/or such a method may be practiced using other structure and/or functionality in addition to or other than one or more of the aspects set forth herein.

[0089] It will also be understood that, although the terms “first,” “second,” etc. may be used herein to describe various elements, these elements should not be limited by these terms. These terms are only used to distinguish one element from another. For example, a first node could be termed a second node, and, similarly, a second node could be termed a first node, which changing the meaning of the description, so long as all occurrences of the “first node” are renamed consistently and all occurrences of the “second node” are renamed consistently. The first node and the second node are both nodes, but they are not the same node.

[0090] The terminology used herein is for the purpose of describing particular implementations only and is not intended to be limiting of the claims. As used in the description of the implementations and the appended claims, the singular forms “a,” “an,” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will also be understood that the term “and/or” as used herein refers to and encompasses any and all possible combinations of one or more of the associated listed items. It will be further understood that the terms “comprises” and/or “comprising,” when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof.

[0091] As used herein, the term “if” may be construed to mean “when” or “upon” or “in response to determining” or “in accordance with a determination” or “in response to detecting,” that a stated condition precedent is true, depending on the context. Similarly, the phrase “if it is determined [that a stated condition precedent is true]” or “if [a stated condition precedent is true]” or “when [a stated condition precedent is true]” may be construed to mean “upon determining” or “in response to determining” or “in accordance with a determination” or “upon detecting” or “in response to detecting” that the stated condition precedent is true, depending on the context.

Claims

1. A method comprising: at a first device with one or more processors and non-transitory memory: determining a deterministic access schedule for a plurality of processes executing on a plurality of devices including the first device and a second device to access a data arrangement, including: calculating a first waking time for a first process of the first device based at least in part at least on a time a first input for the first process is ready to be obtained by the first device; and calculating a second waking time for a second process of the second device, separate from the first device, based at least in part on a time a second input for the second process is ready to be obtained by the second device; and transmitting, to the second device, at least a subset of the deterministic access schedule including the second waking time.
2. The method of claim 1, wherein calculating the second waking time is based on a hardware operation duration.
3. The method of claim 2, wherein the hardware operation duration includes a camera exposure

- time.
4. The method of claim 2, further comprising receiving, from the second device, the hardware operation duration.
 5. The method of claim 1, wherein calculating the second waking time is based on a process execution time of the first process.
 6. The method of claim 1, wherein calculating the second waking time includes adding a reserved communication latency amount of time to the time the second input for the second process is ready to be obtained.
 7. The method of claim 1, wherein calculating the second waking time is based on applying a neural network to one or more process characteristics or hardware parameters of the second device.
 8. The method of claim 7, further comprising receiving, from the second device, the one or more process characteristics or hardware parameters of the second device.
 9. The method of claim 1, further comprising transmitting the second input to the second device.
 10. The method of claim 9, wherein the second input is transmitted to the second device at the time the second input for the second process is ready to be obtained by the second device.
 11. The method of claim 9, wherein the second input is transmitted to the first device at the second waking time.
 12. The method of claim 9, wherein the second input is an image and the second process performs correction on the image prior to display.
 13. The method of claim 12, wherein the first process performs at least one of machine learning, computer vision, or 3D rendering to generate the image.
 14. The method of claim 1, wherein determining the deterministic access schedule further includes calculating a third waking time for a third process of a third device, separate from the first device and the second device, based at least in part on a time a third input for the third process is ready to be obtained by the third device, further comprising transmitting, to the third device, at least a subset of the deterministic access schedule including the third waking time.
 15. A second device comprising: non-transitory memory; and one or more processors to: execute a second process; receive, from a first device separate from the second device, at least a subset of a deterministic access schedule including a second waking time for the second process; and wake the second process at the waking time to access a data arrangement to obtain a second input for the second process.
 16. The second device of claim 15, wherein the second waking time is based on a hardware operation duration and wherein the one or more processors are further to transmit, to the first device, the hardware operation duration.
 17. The second device of claim 15, wherein the second waking time is based on applying a neural network to one or more process characteristics or hardware parameters of the second device and wherein the one or more processors are further to transmit, to the first device, the one or more process characteristics or hardware parameters of the second device.
 18. The device of claim 15, wherein the one or more processors are further to receive the second input from the first device.
 19. The device of claim 18, wherein the second input is an image and the second process performs correction on the image prior to display.
 20. A non-transitory computer-readable medium having instructions encoded thereon which, when executed by one or more processors of a first device, cause the first device to: determine a deterministic access schedule for a plurality of processes executing on a plurality of devices including the first device and a second device to access a data arrangement by: calculating a first waking time for a first process of the first device based at least in part at least on a time a first input for the first process is ready to be obtained by the first device; and calculating a second waking time for a second process of the second device, separate from the first device, based at least in part on a time a second input for the second process is ready to be obtained by the second device; and

transmit, to the second device, at least a subset of the deterministic access schedule including the second waking time.
