

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265193

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

CHACHAD; Abhijeet Ashok et al.

HARDWARE COHERENCE FOR MEMORY CONTROLLER

Abstract

A system includes a non-coherent component; a coherent, non-caching component; a coherent, caching component; and a level two (L2) cache subsystem coupled to the non-coherent component, the coherent, non-caching component, and the coherent, caching component. The L2 cache subsystem includes a L2 cache; a shadow level one (L1) main cache; a shadow L1 victim cache; and a L2 controller. The L2 controller is configured to receive and process a first transaction from the non-coherent component; receive and process a second transaction from the coherent, non-caching component; and receive and process a third transaction from the coherent, caching component.

Inventors: CHACHAD; Abhijeet Ashok (Plano, TX), THOMPSON; David Matthew (Dallas, TX), BHORIA; Naveen (Plano, TX)

Applicant: TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)

Family ID: 1000008576792

Appl. No.: 19/199662

Filed: May 06, 2025

Related U.S. Application Data

parent US continuation 18337521 20230620 parent-grant-document US 12321270 child US 19199662

parent US continuation 17460439 20210830 parent-grant-document US 11687457 child US 18337521

parent US continuation 16882216 20200522 parent-grant-document US 11106584 child US 17460439

us-provisional-application US 62852416 20190524

Publication Classification

Int. Cl.: **G06F12/0811** (20160101); **G06F9/30** (20180101); **G06F9/38** (20180101); **G06F9/46** (20060101); **G06F9/54** (20060101); **G06F11/30** (20060101); **G06F12/0808** (20160101); **G06F12/0815** (20160101); **G06F12/0817** (20160101); **G06F12/0831** (20160101); **G06F12/084** (20160101); **G06F12/0895** (20160101); **G06F12/128** (20160101); **G06F13/16** (20060101)

U.S. Cl.:

CPC **G06F12/0811** (20130101); **G06F9/3004** (20130101); **G06F9/30047** (20130101); **G06F9/30079** (20130101); **G06F9/3867** (20130101); **G06F9/467** (20130101); **G06F9/544** (20130101); **G06F9/546** (20130101); **G06F11/3037** (20130101); **G06F12/0808** (20130101); **G06F12/0815** (20130101); **G06F12/0828** (20130101); **G06F12/0831** (20130101); **G06F12/084** (20130101); **G06F12/0895** (20130101); **G06F12/128** (20130101); **G06F13/1668** (20130101); G06F2212/1021 (20130101); G06F2212/608 (20130101)

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] The present application is a continuation of and claims priority to U.S. patent application Ser. No. 18/337,521, filed Jun. 20, 2023, which claims priority to U.S. patent application Ser. No. 17/460,439, filed Aug. 30, 2021, now U.S. Pat. No. 11,687,457, issued Jun. 27, 2023, which claims priority to U.S. patent application Ser. No. 16/882,216, filed May 22, 2020, now U.S. Pat. No. 11,106,584, issued Aug. 31, 2021, which claims priority to U.S. Provisional Patent Application No. 62/852,416, filed May 24, 2019, each of which is hereby incorporated by reference herein in its entirety.

BACKGROUND

[0002] Some memory systems include a multi-level cache system, in which a hierarchy of memories (e.g., caches) provides varying access speeds to cache data. A first level (L1) cache is closely coupled to a central processing unit (CPU) core and provides the CPU core with relatively fast access to cache data. A second level (L2) cache is also coupled to the CPU core and, in some examples, is larger and thus holds more data than the L1 cache, although the L2 cache provides relatively slower access to cache data than the L1 cache. Additional memory levels of the hierarchy are possible.

SUMMARY

[0003] In accordance with at least one example of the disclosure, a method includes determining, by a level one (L1) controller, to change a size of a L1 main cache; servicing, by the L1 controller, pending read requests and pending write requests from a central processing unit (CPU) core; stalling, by the L1 controller, new read requests and new write requests from the CPU core; writing back and invalidating, by the L1 controller, the L1 main cache. The method also includes receiving, by a level two (L2) controller, an indication that the L1 main cache has been invalidated and, in response, flushing a pipeline of the L2 controller; in response to the pipeline being flushed, stalling, by the L2 controller, requests received from any master; reinitializing, by the L2 controller, a shadow L1 main cache. Reinitializing includes clearing previous contents of the shadow L1 main cache and changing the size of the shadow L1 main cache.

[0004] In accordance with at least one example of the disclosure, an apparatus, includes a central processing unit (CPU) core and a level one (L1) cache subsystem coupled to the CPU core. The L1

cache subsystem includes a L1 main cache and a L1 controller. The L1 controller is configured to determine to change a size of the L1 main cache, service pending read requests and pending write requests from the CPU core, stall new read requests and new write requests from the CPU core, and write back and invalidate the L1 main cache. The apparatus also includes a level two (L2) cache subsystem coupled to the L1 cache subsystem. The L2 cache subsystem includes a L2 main cache, a shadow L1 main cache, and a L2 controller. The L2 controller is configured to receive an indication that the L1 main cache has been invalidated and, in response, flush a pipeline of the L2 controller; in response to the pipeline being flushed, stall requests received from any master; and reinitialize the shadow L1 main cache. Reinitializing includes clearing previous contents of the shadow L1 main cache and changing the size of the shadow L1 main cache.

[0005] In accordance with another example of the disclosure, a method includes receiving, by a level two (L2) controller, an indication that the a level one (L1) main cache has been invalidated and, in response, flushing a pipeline of the L2 controller; in response to the pipeline being flushed, stalling, by the L2 controller, requests received from any master; and reinitializing, by the L2 controller, a shadow L1 main cache. Reinitializing includes clearing previous contents of the shadow L1 main cache and changing the size of the shadow L1 main cache.

[0006] In accordance with at least one example of the disclosure, an apparatus includes a CPU core and a L1 cache subsystem coupled to the CPU core. The L1 cache subsystem includes a L1 main cache, a L1 victim cache, and a L1 controller. The apparatus includes a L2 cache subsystem coupled to the L1 cache subsystem. The L2 cache subsystem includes a L2 main cache, a shadow L1 main cache, a shadow L1 victim cache, and a L2 controller. The L2 controller receives an indication from the L1 controller that a cache line A is being relocated from the L1 main cache to the L1 victim cache; in response to the indication, update the shadow L1 main cache to reflect that the cache line A is no longer located in the L1 main cache; and in response to the indication, update the shadow L1 victim cache to reflect that the cache line A is located in the L1 victim cache.

[0007] In accordance with at least one example of the disclosure, a method includes receiving, by a level two (L2) controller of a L2 cache subsystem, an indication from a level one (L1) cache subsystem that a cache line A is being relocated from a L1 main cache to a L1 victim cache; in response to the indication, updating, by the L2 controller, a shadow L1 main cache of the L2 cache subsystem to reflect that the cache line A is no longer located in the L1 main cache; and in response to the indication, updating a shadow L1 victim cache of the L2 cache subsystem to reflect that the cache line A is located in the L1 victim cache.

[0008] In accordance with at least one example of the disclosure, a level two (L2) cache subsystem includes a L2 main cache; a shadow level one (L1) main cache; a shadow L1 victim cache; and a L2 controller. The L2 controller is configured to: receive an indication from a level one (L1) controller that a cache line A is being relocated from a L1 main cache to a L1 victim cache; in response to the indication, update the shadow L1 main cache to reflect that the cache line A is no longer located in the L1 main cache; and in response to the indication, update the shadow L1 victim cache to reflect that the cache line A is located in the L1 victim cache.

[0009] In accordance with at least one example of the disclosure, a system includes a non-coherent component; a coherent, non-caching component; a coherent, caching component; and a level two (L2) cache subsystem coupled to the non-coherent component, the coherent, non-caching component, and the coherent, caching component. The L2 cache subsystem includes a L2 cache; a shadow level one (L1) main cache; a shadow L1 victim cache; and a L2 controller. The L2 controller is configured to receive and process a first transaction from the non-coherent component; receive and process a second transaction from the coherent, non-caching component; and receive and process a third transaction from the coherent, caching component.

[0010] In accordance with another example of the disclosure, a method includes maintaining, by a level two (L2) cache controller, a L2 cache, a shadow level one (L1) main cache and a shadow L1 victim cache; receiving and processing, by the L2 cache controller, a first transaction from a non-

coherent component; receiving and processing, by the L2 cache controller, a second transaction from a coherent, non-caching component; and receiving and processing, by the L2 cache controller, a third transaction from a coherent, caching component.

[0011] In accordance with at least one example of the disclosure, an apparatus includes a CPU core and a L1 cache subsystem including a L1 main cache, a L1 victim cache, and a L1 controller. The apparatus includes a L2 cache subsystem including a L2 main cache, a shadow L1 main cache, a shadow L1 victim cache, and a L2 controller configured to receive a read request from the L1 controller as a single transaction. Read request includes a read address, a first indication of an address and a coherence state of a cache line A to be moved from the L1 main cache to the L1 victim cache to allocate space for data returned in response to the read request, and a second indication of an address and a coherence state of a cache line B to be removed from the L1 victim cache in response to the cache line A being moved to the L1 victim cache.

[0012] In accordance with at least one example of the disclosure, a method includes receiving, by a level two (L2) controller, a read request from a level one (L1) controller as a single transaction. The read request includes a read address, a first indication of an address and a coherence state of a cache line A to be moved from a L1 main cache to a L1 victim cache to allocate space for data returned in response to the read request, and a second indication of an address and a coherence state of a cache line B to be removed from the L1 victim cache in response to the cache line A being moved to the L1 victim cache.

[0013] In accordance with at least one example of the disclosure, a level two (L2) cache subsystem, includes a L2 main cache, a shadow L1 main cache, a shadow L1 victim cache, and a L2 controller configured to receive a read request in a single transaction from a level one (L1) controller. The read request includes a read address, a first indication of an address and a coherence state of a cache line A to be moved from a L1 main cache to a L1 victim cache to allocate space for data returned in response to the read request, and a second indication of an address and a coherence state of a cache line B to be removed from the L1 victim cache in response to the cache line A being moved to the L1 victim cache.

[0014] In accordance with at least one example of the disclosure, an apparatus includes a CPU core and a L1 cache subsystem including a L1 main cache, a L1 victim cache, and a L1 controller. The apparatus includes a L2 cache subsystem coupled to the L1 cache subsystem by a transaction bus and a tag update bus. The L2 cache subsystem includes a L2 main cache, a shadow L1 main cache, a shadow L1 victim cache, and a L2 controller. The L2 controller receives a message from the L1 controller over the tag update bus, including a valid signal, an address, and a coherence state. In response to the valid signal being asserted, the L2 controller identifies an entry in the shadow L1 main cache or the shadow L1 victim cache having an address corresponding to the address of the message and updates a coherence state of the identified entry to be the coherence state of the message.

[0015] In accordance with at least one example of the disclosure, a method includes receiving, by a level two (L2) controller, a message from a level one (L1) controller over a tag update bus separate from a transaction bus between the L2 controller and the L1 controller, the message comprising a valid signal, an address, and a coherence state. The method also includes, in response to the valid signal being asserted, identifying, by the L2 controller, an entry in a shadow L1 main cache or a shadow L1 victim cache having an address corresponding to the address of the message; and updating, by the L2 controller, a coherence state of the identified entry to be the coherence state of the message.

[0016] In accordance with at least one example of the disclosure, an apparatus includes a central processing unit (CPU) core and a level one (L1) cache subsystem coupled to the CPU core. The L1 cache subsystem includes a L1 main cache, a L1 victim cache, and a L1 controller. The apparatus includes a level two (L2) cache subsystem coupled to the L1 cache subsystem by a transaction bus and a tag update bus. The L2 cache subsystem includes a L2 main cache, a shadow L1 main cache,

a shadow L1 victim cache, and a L2 controller. The L2 controller is configured to receive a message from the L1 controller over the tag update bus, the message comprising a valid signal and an address. In response to the valid signal being asserted, the L2 controller identifies an entry in the shadow L1 victim cache having an address corresponding to the address of the message and updates a coherence state of the identified entry to be invalid.

[0017] In accordance with at least one example of the disclosure, a method includes receiving, by a level two (L2) controller, a first request for a cache line in a shared cache coherence state; mapping, by the L2 controller, the first request to a second request for a cache line in an exclusive cache coherence state; and responding, by the L2 controller, to the second request.

[0018] In accordance with at least one example of the disclosure, a method includes receiving, by a level two (L2) controller, a first request for a cache line in a shared cache coherence state; determining, by the L2 controller, that the cache line is not present in a L2 cache; mapping, by the L2 controller, the first request to a second request for a cache line in an exclusive cache coherence state; and forwarding, by the L2 controller, the second request.

[0019] In accordance with at least one example of the disclosure, an apparatus includes a central processing unit (CPU) core, a level one (L1) cache subsystem coupled to the CPU core, and a level two (L2) cache subsystem coupled to the L1 cache subsystem. The L2 cache subsystem includes a L2 main cache and a L2 controller. The L2 controller is configured to receive a first request for a cache line in a shared cache coherence state, map the first request to a second request for a cache line in an exclusive cache coherence state, and respond to the second request.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0020] For a detailed description of various examples, reference will now be made to the accompanying drawings in which:

[0021] FIG. 1 shows a block diagram of a multi-level cache system in accordance with various examples;

[0022] FIG. 2 shows another block diagram of a multi-level cache system in accordance with various examples;

[0023] FIG. 3 shows another block diagram of a multi-level cache system including level one (L1) main and victim caches in accordance with various examples;

[0024] FIGS. 4a-4f show flow charts of methods for processing various transaction types in accordance with various examples;

[0025] FIG. 5 shows an example organization of the L1 main cache and the L1 victim cache in accordance with various examples;

[0026] FIG. 6 shows an example organization of level two (L2) shadow L1 main cache and shadow L1 victim cache in accordance with various examples;

[0027] FIG. 7 shows an example of a read allocate in the L1 main and victim caches in accordance with various examples;

[0028] FIG. 8 shows an example of a read allocate in the L2 shadow L1 main and shadow L1 victim caches in accordance with various examples;

[0029] FIG. 9 shows a table of sideband signaling protocol data in accordance with various examples;

[0030] FIG. 10 shows a table of tag update protocol data in accordance with various examples;

[0031] FIG. 11 shows a block diagram illustrating shared state suppression in accordance with various examples;

[0032] FIG. 12 shows a flow chart of a method for shared state suppression in accordance with various examples; and

[0033] FIG. 13 shows a flow chart of a method for changing a cache size in accordance with various examples.

DETAILED DESCRIPTION

[0034] FIG. 1 shows a block diagram of a system **100** in accordance with an example of this disclosure. The example system **100** includes multiple CPU cores **102a-102n**. Each CPU core **102a-102n** is coupled to a dedicated L1 cache **104a-104n** and a dedicated L2 cache **106a-106n**. The L2 caches **106a-106n** are, in turn, coupled to a shared third level (L3) cache **108** and a shared main memory **110** (e.g., double data rate (DDR) random-access memory (RAM)). In other examples, a single CPU core **102** is coupled to a L1 cache **104**, a L2 cache **106**, a L3 cache **108**, and main memory **110**.

[0035] In some examples, the CPU cores **102a-102n** include a register file, an integer arithmetic logic unit, an integer multiplier, and program flow control units. In an example, the L1 caches **104a-104n** associated with each CPU core **102a-102n** include a separate level one program cache (L1P) and level one data cache (L1D). The L2 caches **106a-106n** are combined instruction/data caches that hold both instructions and data. In certain examples, a CPU core **102a** and its associated L1 cache **104a** and L2 cache **106a** are formed on a single integrated circuit.

[0036] The CPU cores **102a-102n** operate under program control to perform data processing operations upon data. Instructions are fetched before decoding and execution. In the example of FIG. 1, L1P of the L1 cache **104a-104n** stores instructions used by the CPU cores **102a-102n**. A CPU core **102** first attempts to access any instruction from L1P of the L1 cache **104**. L1D of the L1 cache **104** stores data used by the CPU core **102**. The CPU core **102** first attempts to access any required data from L1 cache **104**. The two L1 caches **104** (L1P and L1D) are backed by the L2 cache **106**, which is a unified cache. In the event of a cache miss to the L1 cache **104**, the requested instruction or data is sought from L2 cache **106**. If the requested instruction or data is stored in the L2 cache **106**, then it is supplied to the requesting L1 cache **104** for supply to the CPU core **102**. The requested instruction or data is simultaneously supplied to both the requesting cache and CPU core **102** to speed use.

[0037] The unified L2 cache **106** is further coupled to a third level (L3) cache **108**, which is shared by the L2 caches **106a-106n** in the example of FIG. 1. The L3 cache **108** is in turn coupled to a main memory **110**. As will be explained in further detail below, memory controllers facilitate communication between various ones of the CPU cores **102**, the L1 caches **104**, the L2 caches **106**, the L3 cache **108**, and the main memory **110**. The memory controller(s) handle memory centric functions such as cacheability determination, cache coherency implementation, error detection and correction, address translation and the like. In the example of FIG. 1, the CPU cores **102** are part of a multiprocessor system, and thus the memory controllers also handle data transfer between CPU cores **102** and maintain cache coherence among CPU cores **102**. In other examples, the system **100** includes only a single CPU core **102** along with its associated L1 cache **104** and L2 cache **106**.

[0038] FIG. 2 shows a block diagram of a system **200** in accordance with examples of this disclosure. Certain elements of the system **200** are similar to those described above with respect to FIG. 1, although shown in greater detail. For example, a CPU core **202** is similar to the CPU core **102** described above. The L1 cache **104** subsystem described above is depicted as L1D **204** and L1P **205**. The L2 cache **106** described above is shown here as L2 cache subsystem **206**. An L3 cache **208** is similar to the L3 cache **108** described above. The system **200** also includes a streaming engine **210** coupled to the L2 cache subsystem **206**. The system **200** also includes a memory management unit (MMU) **207** coupled to the L2 cache subsystem **206**.

[0039] The L2 cache subsystem **206** includes L2 tag ram **212**, L2 coherence (e.g., MESI) data **214**, shadow L1 tag ram **216**, and L1 coherence (e.g., MESI) data **218**. Each of the blocks **212**, **214**, **216**, **218** are alternately referred to as a memory or a RAM. The L2 cache subsystem **206** also includes tag ram error correcting code (ECC) data **220**. In an example, the ECC data **220** is maintained for each of the memories **212**, **214**, **216**, **218**.

[0040] The L2 cache subsystem **206** includes L2 controller **222**, the functionality of which will be described in further detail below. In the example of FIG. 2, the L2 cache subsystem **206** is coupled to memory (e.g., L2 SRAM **224**) including four banks **224a-224d**. An interface **230** performs data arbitration functions and generally coordinates data transmission between the L2 cache subsystem **206** and the L2 SRAM **224**, while an ECC block **226** performs error correction functions. The L2 cache subsystem **206** includes one or more control or configuration registers **228**.

[0041] In the example of FIG. 2, the L2 SRAM is depicted as four banks **224a-224d**. However, in other examples, the L2 SRAM includes more or fewer banks, including being implemented as a single bank. The L2 SRAM **224** serves as the L2 cache and is alternately referred to herein as L2 cache **224**.

[0042] The L2 tag ram **212** includes a list of the physical addresses whose contents (e.g., data or program instructions) have been cached to the L2 cache **224**. In an example, an address translator translates virtual addresses to physical addresses. In one example, the address translator generates the physical address directly from the virtual address. For example, the lower n bits of the virtual address are used as the least significant n bits of the physical address, with the most significant bits of the physical address (above the lower n bits) being generated based on a set of tables configured in main memory. In this example, the L2 cache **224** is addressable using physical addresses. In certain examples, a hit/miss indicator from a tag ram **212** look-up is stored.

[0043] The L2 MESI memory **214** maintains coherence data to implement full MESI coherence with L2 SRAM **224**, external shared memories, and data cached in L2 cache from other places in the system **200**. The functionalities of system **200** coherence are explained in further detail below.

[0044] The L2 cache subsystem **206** also shadows L1D tags in the L1D shadow tag ram **216** and L1D MESI memory **218**. The tag ram ECC data **220** provides error detection and correction for the tag memories and, additionally, for one or both of the L2 MESI memory **214** and the L1D MESI memory **218**. The L2 cache controller **222** generally controls the operations of the L2 cache subsystem **206**, including handling coherency operations both internal to the L2 cache subsystem **206** and among the other components of the system **200**.

[0045] FIG. 3 shows a block diagram of a system **300** that demonstrates various features of cache coherence implemented in accordance with examples of this disclosure. The system **300** contains elements similar to those described above with respect to FIGS. 1 and 2. For example, the CPU core **302** is similar to the CPU cores **102**, **202**. FIG. 3 also includes a L1 cache subsystem **304**, a L2 cache subsystem **306**, and an L3 cache subsystem **308**. The L1 cache subsystem **304** includes a L1 controller **310** coupled to L1 SRAM **312**. The L1 controller **310** is also coupled to a L1 main cache **314** and a L1 victim cache **316**, which are explained in further detail below. In some examples, the L1 main and victim caches **314**, **316** implement the functionality of L1D **204** and/or L1P **205**.

[0046] The L1 controller **310** is coupled to a L2 controller **320** of the L2 cache subsystem **306**. The L2 controller **320** also couples to L2 SRAM **322**. The L2 controller **320** couples to a L2 cache **324** and to a shadow of the L1 main cache **326** as well as a shadow of the L1 victim cache **328**. L2 cache **324** and L2 SRAM **322** are shown separately for ease of discussion, although may be implemented physically together (e.g., as part of L2 SRAM **224**, including in a banked configuration, as described above. Similarly, the shadow L1 main cache **326** and the shadow L1 victim cache **328** may be implemented physically together, and are similar to the L1D shadow tag ram **216** and the L1D MESI **218**, described above. The L2 controller **320** is also coupled to a L3 controller **309** of the L3 cache subsystem **308**. L3 cache and main memory (e.g., DDR **110** described above) are not shown for simplicity.

[0047] Cache coherence is a technique that allows data and program caches, as well as different requestors (including requestors that do not have caches) to determine the most current data value for a given address in memory. Cache coherence enables this coherent data value to be accurately reflected to observers (e.g., a cache or requestor that issues commands to read a given memory location) present in the system **300**. Certain examples of this disclosure refer to an exemplary

MESI coherence scheme, in which a cache line is set to one of four cache coherence states: modified, exclusive, shared, or invalid. Other examples of this disclosure refer to a subset of the MESI coherence scheme, while still other examples include more coherence states than the MESI coherence scheme.

[0048] Regardless of the coherence scheme, cache coherence states for a given cache line are stored in, for example, the L2 MESI memory **214** described above.

[0049] A cache line having a cache coherence state of modified indicates that the cache line is modified with respect to main memory (e.g., DDR **110**), and the cache line is held exclusively in the current cache (e.g., the L2 cache **324**). A modified cache coherence state also indicates that the cache line is explicitly not present in any other caches (e.g., L1 or L3 caches).

[0050] A cache line having a cache coherence state of exclusive indicates that the cache line is not modified with respect to main memory (e.g., DDR **110**), but the cache line is held exclusively in the current cache (e.g., the L2 cache **324**). An exclusive cache coherence state also indicates that the cache line is explicitly not present in any other caches (e.g., L1 or L3 caches).

[0051] A cache line having a cache coherence state of shared indicates that the cache line is not modified with respect to main memory (e.g., DDR **110**). A shared cache state also indicates that the cache line may be present in multiple caches (e.g., caches in addition to the L2 cache **324**).

[0052] A cache line having a cache coherence state of invalid indicates that the cache line is not present in the cache (e.g., the L2 cache **324**).

[0053] Examples of this disclosure leverage hardware techniques, control logic, and/or state information to implement a coherent system. Each observer can issue read requests—and certain observers are able to issue write requests—to memory locations that are marked shareable. Caches in particular can also have snoop requests issued to them, requiring their cache state to be read, returned, or even updated, depending on the type of the snoop operation. In the exemplary multi-level cache hierarchy described above, the L2 cache subsystem **306** is configured to both send and receive snoop operations. The L1 cache subsystem **304** receives snoop operations, but does not send snoop operations. The L3 cache subsystem **308** sends snoop operations, but does not receive snoop operations. In examples of this disclosure, the L2 cache controller **320** maintains state information (e.g., in the form of hardware buffers, memories, and logic) to additionally track the state of coherent cache lines present in both the L1 main cache **314** and the L1 victim cache **316**. Tracking the state of coherent cache lines enables the implementation of a coherent hardware cache system.

[0054] Examples of this disclosure refer to various types of coherent transactions, including read transactions, write transactions, snoop transactions, victim transactions, and cache maintenance operations (CMO). These transactions are at times referred to as reads, writes, snoops, victims, and CMOs, respectively.

[0055] Reads return the current value for a given address, whether that value is stored at the endpoint (e.g., DDR **110**), or in one of the caches in the coherent system **300**. Writes update the current value for a given address, and invalidate other copies for the given address stored in caches in the coherent system **300**. Snoops read or invalidate (or both) copies of data stored in caches. Snoops are initiated from a numerically-higher level of the hierarchy to a cache at the next, numerically-lower level of the hierarchy (e.g., from the L2 controller **320** to the L1 controller **310**), and are able to be further propagated to even lower levels of the hierarchy as needed. Victims are initiated from a numerically-lower level cache in the hierarchy to the next, numerically-higher level of the cache hierarchy (e.g., from the L1 controller **310** to the L2 controller **320**). Victims transfer modified data to the next level of the hierarchy. In some cases, victims are further propagated to numerically-higher levels of the cache hierarchy (e.g., if the L2 controller **310** sends a victim to the L2 controller **320** for an address in the DDR **110**, and the line is not present in the L2 cache **324**, the L2 controller **320** forwards the victim to the L3 controller **309**). Finally, CMOs cause an action to be taken in one of the caches for a given address.

[0056] Still referring to FIG. 3, in one example, the L1 main cache **314** is a direct mapped cache that services read and write hits and snoops. The L1 main cache **314** also keeps track of cache coherence state information (e.g., MESI state) for its cache lines. In an example, the L1 main cache **314** is a read-allocate cache. Thus, writes that miss the L1 main cache **314** are sent to L2 cache subsystem **306** without allocating space in the L1 main cache **314**. In the example where the L1 main cache **314** is direct mapped, when a new allocation takes place in the L1 main cache **314**, the current line in the set is moved to the L1 victim cache **316**, regardless of whether the line is clean (e.g., unmodified) or dirty (e.g., modified).

[0057] In an example, the L1 victim cache **316** is a fully associative cache that holds cache lines that have been removed from the L1 main cache **314**, for example due to replacement. The L1 victim cache **316** holds both clean and dirty lines. The L1 victim cache **316** services read and write hits and snoops. The L1 victim cache **316** also keeps track of cache coherence state information (e.g., MESI state) for its cache lines. When a cache line in the modified state is replaced from the L1 victim cache **316**, that cache line is sent to the L2 cache subsystem **306** as a victim.

[0058] As explained above, the L2 cache subsystem **306** includes a unified L2 cache **324** that is used to service requests from multiple requestor types, including L1D and L1P (through the L1 controller **310**), the streaming engine **210**, a memory management unit (MMU **207**), and the L3 cache (through the L3 controller **309**). In an example, the L2 cache **324** is non-inclusive with the L1 cache subsystem **304**, which means that the L2 cache **324** is not required to include all cache lines stored in the L1 caches **314**, **316**, but that some lines may be cached in both levels.

Continuing this example, the L2 cache **324** is also non-exclusive, which means that cache lines are not explicitly prevented from being cached in both the L1 and L2 caches **314**, **316**, **324**. For example, due to allocation and random replacement, cache lines may be present in one, both, or neither of the L1 and L2 caches. The combination of non-inclusive and non-exclusive cache policies enables the L2 controller **320** to manage its cache contents without requiring the L1 controller **310** to invalidate or remove cache lines. This simplifies processing in the L2 cache subsystem **306** and enables increased performance for the CPU core **302** by allowing critical data to remain cached in the L1 cache subsystem **304** even if it has been evicted from the L2 cache **324**.

[0059] Still referring to FIG. 3, the L2 controller **320** described herein combines both local coherence (e.g., handling requests targeting its local L2 SRAM **322** as an endpoint) and external coherence (e.g., handling requests targeting external memories, such as L3 SRAM (not shown for simplicity) or DDR **110** as endpoints). An endpoint refers to a memory target such as L2 SRAM **322** or DDR **110** that resides at a particular location on the chip, is acted upon directly by a single controller and/or interface, and may be cached at various levels of a coherent cache hierarchy, such as depicted in FIG. 3. A master (e.g., a hardware component, circuitry, or the like) refers to a requestor that issues read and write accesses to an endpoint. In some examples, a master stores the results of these read and write accesses in a cache, although the master does not necessarily store such results in a cache.

[0060] Local coherence requests are received by the L2 controller **320** from, for example, the CPU core **302** or as a direct memory access (DMA) request from another CPU core or a master associated with another CPU core. External coherence requests are received by the L2 controller **320** from, for example, the CPU core **302** or L3 controller **309**. Thus, the single L2 controller **320** is configured to address both local and external coherence.

[0061] In accordance with various examples, the L2 controller **320** manages the CPU core **302** coherent view of three endpoints: L2 SRAM **322**, L3 SRAM (part of the L3 cache subsystem **308**, not shown for simplicity), and main memory or DDR **110**, described above. For ease of discussion, L3 SRAM and DDR **110** are grouped together and referred to as an “external” memory or endpoint, which distinguishes them from the L2 SRAM **322** as a “local” (e.g., to the L2 controller **320**) memory or endpoint.

[0062] A master refers to a requestor that issues read and write accesses to an endpoint. In some

examples, a master stores the results of these read and write accesses in a cache, although the master does not necessarily store such results in a cache. Coherent masters (e.g., masters for whom coherence must be handled by L2 controller 320) are classified as either caching or non-caching. Non-coherent masters (e.g., masters that do not require coherent data) are not distinguished as caching or non-caching due to their being non-coherent. Referring briefly back to FIG. 2, in some examples non-coherent masters include L1P 205. Coherent, non-caching masters include MMU 207, SE 210, and L3 208. Coherent, caching masters include L1D 204.

[0063] The L2 controller 320 is configured to provide coherent access to both internal and external endpoints for coherent masters, while also providing access to those internal and external endpoints for non-coherent masters. As will be explained in further detail below, the L2 controller manages coherent state information, issues coherence transactions (e.g., snoop, victim) to maintain proper coherence states, and propagates information as needed to the downstream controllers such as the L3 controller 309 to provide a coherent view of the memory stored in the L2 cache subsystem 306.

[0064] As will be explained further below, the L2 controller 320 is configured to perform normal cache allocation, replacement, and victimization operations, while also sending coherent transactions to communicate the storage of coherent locations within the L2 cache subsystem 306 or L1 cache subsystem 304. As a result, downstream cache controllers such as the L3 controller 309 are able to maintain the directory information, if so enabled, about what addresses are held in the L1 and L2 cache subsystems 304, 306.

[0065] In accordance with examples of this disclosure, the L2 controller 320 is part of a system that includes a non-coherent master; a non-caching, coherent master; and a caching, coherent master. The L2 controller 320 is configured to receive and process transactions from each of these masters, while maintaining global coherence (e.g., with respect to external memories) and local coherence (e.g., with respect to its local memory) as required by the particular master. Thus, the L2 controller 320 also enables interleaving of coherent and non-coherent traffic among the various masters.

[0066] The following table summarizes interactions between various masters and the L2 controller 320 in accordance with various examples. In particular, Table 1 indicates for a particular master what transaction types that master can initiate to the L2 controller 320, what transaction types the L2 controller 320 can initiate to that master, and whether global and/or local coherence is supported by the L2 controller 320 for that master.

TABLE-US-00001 TABLE 1 Master-initiated L2-initiated Global Local Master transaction									
transaction		coherence?		coherence?		L1P 205			
Read		None		No		Yes			
MMU 207		Read		None		Yes			
SE 210		Read, CMO		None		Yes			
L1D 204		R, W, Victim		Snoop		Yes			
L3 208		Snoop		R, W, Victim		Yes			
DMA		Read, Write		None		No			

[0067] FIG. 4a shows a method 400 carried out by the L2 controller 320 in response to a read request from a non-coherent master, such as L1P 205. The method 400 begins in block 402 with the L2 controller 320 receiving a read request from a non-coherent master, and continues in block 404 with reading data from an endpoint based on the read request. Although not explicitly shown, if the read request hits in the L2 cache 324, the L2 controller 320 is configured to read the data from the L2 cache 324. On the other hand, if the read request does not hit in the L2 cache 324, the L2 controller 320 is configured to read the data from an endpoint, such as the L3 cache subsystem 308 or DDR 110. Once the L2 controller 320 has read response data (either from L2 cache 324 or from an endpoint), the method 400 continues to block 406 in which the L2 controller 320 returns the read response data to the non-coherent master.

[0068] FIG. 4b shows a method 410 carried out by the L2 controller 320 in response to a read request from a coherent, non-caching master, such as MMU 207, SE 210, and L3 208. When coherent, non-caching masters issue read commands to the L2 controller 320, either to a local endpoint or external endpoint, the L2 controller 320 determines if the line is present in the L1 cache 314, 316, and if so, whether a snoop command should be issued to obtain the latest copy from L1 caches 314, 316, or if the data can be obtained from the endpoint (e.g., L2 SRAM 322) or the L2

cache **324** (if present). Due to variations in access latency for a local endpoint (faster) compared to an external endpoint (slower), the L2 controller **320** makes multiple decisions for where and how to obtain a coherent memory location in response to a read command from a non-caching master. [0069] The method **410** begins in block **412** with the L2 controller **320** receiving a read request from a coherent, non-caching master, and continues in block **414** with the L2 controller **320** determining whether the read request hits in the shadow L1 main cache **326** or the shadow L1 victim cache **328**, which indicates that the requested data may be present in the L1 cache subsystem **304**.

[0070] If, in block **414**, the read request does not hit in the shadow L1 caches **326**, **328**, the method **410** continues in block **416** in which the L2 controller **320** reads the data from an endpoint and returns the data as a read response. However, if in block **414** the read request hits one of the shadow L1 caches **326**, **328**, the method **410** continues in block **418** with the L2 controller **320** generating a snoop read to the L1 controller **310**. If a snoop response from the L1 controller **310** contains valid data in block **420**, then the L2 controller **320** returns the snoop response data as a read response to the requesting master in block **422**. If the snoop response from the L1 controller **310** contains invalid data in block **420**, then the L2 controller **320** returns endpoint data as the read response to the requesting master in block **416**.

[0071] FIG. **4c** shows a method **430** carried out by the L2 controller **320** in response to a read request from a coherent, caching master, such as L1D **204**. The method **430** makes reference to various sideband signals that describe allocations that will occur in the L1 cache subsystem **304** (e.g., movements of cache lines in L1 main cache **314** and L1 victim cache **316**) as a result of the read request. These sideband signals are described in further detail below with respect to FIG. **9**. [0072] In particular, the method **430** begins in block **432** with the L2 controller **320** receiving an allocating read request from a coherent, caching master, which in this example is the L1 cache subsystem **304**. This read request includes sideband signals that indicate it is an allocating request (e.g., alloc==1). The method **430** then proceeds to block **434** in which the L2 controller **320** writes an address and, optionally, a secure bit, indicated by the sideband signals to the shadow L1 main cache **326**, which now indicates the address that is being allocated to the L1 main cache **314** as a result of this read request.

[0073] The method **430** continues in block **436** with determining whether a main_valid sideband signal is asserted, which indicates that a cache line is moving from the L1 main cache **314** to the L1 victim cache **316** as a result of this read request. If the main_valid signal is asserted, the method **430** continues to block **438** in which the L2 controller **320** updates its shadow L1 victim cache **328** to include an address specified by main_address, a coherence state specified by main_mesi, and optionally a secure bit specified by main_secure. As a result, the shadow L1 victim cache **328** now includes the address and coherence state information of the line that is being moved from the L1 main cache **314** to the L1 victim cache **316** as a result of this read request.

[0074] If the main_valid signal is de-asserted, then a line is not being moved from the L1 main cache **314** to the L1 victim cache **316** as a result of this read request, and the method **430** continues to block **440** with determining whether a victim_valid sideband signal is asserted, which indicates that a cache line is moving out of the L1 victim cache **316** as a result of this read request (e.g., is being displaced by the L1 main cache **314** to L1 victim cache **316** movement described above). If the victim_valid signal is asserted, the method **430** continues in block **442** with determining whether the coherence state specified by victim_mesi (e.g., the coherence state of the line being moved out of the L1 victim cache **316**) is invalid, modified, or shared/exclusive.

[0075] If victim_mesi is invalid, the method **430** proceeds to block **448** in which the L2 controller **320** returns read response data from an endpoint, or the L2 cache **324**.

[0076] If victim_mesi is shared/exclusive, the method **430** continues to block **444** where the L2 controller **320** removes an entry from its shadow L1 victim cache **328** having an address that matches victim_address and, optionally, victim_secure. As explained further below, the L2

controller **320** removes the entry in this case because a subsequent victim transaction from the L1 controller **310** does not result when the line evicted from L1 victim cache **316** is in the shared/exclusive state, and thus it is safe to also remove from the shadow L1 victim cache **328**. The method **430** then proceeds to block **448** in which the L2 controller **320** returns read response data from an endpoint, or the L2 cache **324**.

[0077] If victim_mesi is modified, the method **430** continues to block **446** where the L2 controller **320** retains an entry from its shadow L1 victim cache **328** having an address that matches victim_address and, optionally, victim_secure. As explained further below, the L2 controller **320** retains the entry in this case because a subsequent victim transaction from the L1 controller **310** is expected when the line evicted from L1 victim cache **316** is in the modified state. The method **430** then proceeds to block **448** in which the L2 controller **320** returns read response data from an endpoint, or the L2 cache **324**.

[0078] FIG. **4d** shows a method **450** carried out by the L2 controller **320** in response to a write request from a coherent, non-caching master, such as a DMA request from a different CPU core. The method **450** begins in block **452** when the write request is received and continues to block **454** with the L2 controller **320** determining whether the write request hits in the shadow L1 main or victim caches **326**, **328**. If the write request does not hit in the shadow L1 main or victim caches **326**, **328**, then the L2 controller **320** does not need to invalidate any line in the L1 cache subsystem **304** and the method **450** proceeds to block **456** where the L2 controller **320** writes the data to an endpoint.

[0079] However, if the write request hits in the shadow L1 main or victim caches **326**, **328**, then the method **450** proceeds to block **458** in which the L2 controller **320** issues a snoop read and invalidate request to the L1 cache subsystem **304**. If the snoop response has dirty (e.g., modified) data in block **460**, then the L2 controller **320** merges the write data over the snoop response data and writes to an endpoint in block **462**. If the snoop response contains unmodified data in block **460**, then the L2 controller **320** writes the write data to the endpoint in block **456**.

[0080] FIG. **4e** shows a method **470** carried out by the L2 controller **320** in response to a victim from a L1D **204**, which is a coherent, caching master. The method **470** begins in block **472** with the L2 controller **320** receiving a victim from the L1 controller **310**. If a victim address and, optionally, secure bit hits in the shadow L1 victim cache **328** in block **474**, the L2 controller **320** is configured to update the shadow L1 victim cache **328** to invalidate a corresponding address if necessary. The method **470** then continues in block **478**, in which the L2 controller **320** updates an endpoint with the victim data. However, if the victim address and, optionally, secure bit does not hit in the shadow L1 victim cache **328** in block **474**, then the method **470** proceeds to block **478** and the L2 controller **320** updates an endpoint with the victim data without modifying the shadow L1 victim cache **328**.

[0081] FIG. **4f** shows a method **480** carried out by the L2 controller **320** in response to a snoop command from L3 **208**, which is a coherent, non-caching master. The method **480** begins in block **482** in which the L2 controller **320** receives a snoop request from the L3 controller **309**. If, in block **484**, the snoop request hits in the shadow L1 main or victim caches **326**, **328**, the method **480** continues in block **486** with the L2 controller **320** issuing a snoop read **486** to the L1 controller **310**. The method **480** then continues in block **488** with the L2 controller **320** determining whether the snoop response from the L1 controller **310** has valid data.

[0082] If the snoop response from the L1 controller **310** contains invalid data (or if the snoop request did not hit in the shadow L1 main or victim caches **326**, **328** in block **484**), the method **480** continues to block **490** in which the L2 controller **320** determines whether the snoop read hits in the L2 cache **324**. If the snoop read does not hit in the L2 cache **324**, the method **480** continues to block **492** and the L2 controller **320** issues a snoop miss to the L3 controller **309**. However, if the snoop read hits in the L2 cache **324**, the method **480** continues to block **493** in which the L2 controller **320** reads the data from the L2 cache **324** and to block **494** in which the L2 controller

320 updates a coherence state as needed. Then the L2 controller **320** returns the data from the L2 cache **324** as snoop response data to the L3 controller **309** in block **495**.

[0083] If the snoop response from the L1 controller **310** contains valid data in block **488**, the method **480** continues to block **496** in which the L2 controller **320** determines whether the snoop response from the L1 controller **310** hits in the L2 cache **324**. If the snoop response from the L1 controller **310** hits in the L2 cache **324**, the method **480** continues to block **497** in which the L2 controller **320** updates a coherence state of the L2 cache **324** as needed. Then, the L2 controller **320** returns the snoop response data from the L1 controller **310** as a snoop response to the L3 controller **309** in block **498**. If the snoop response from the L1 controller **310** does not hit in the L2 cache **324**, the method **480** proceeds directly to block **498**, in which the L2 controller **320** returns the snoop response data from the L1 controller **310** as a snoop response to the L3 controller **309**.

[0084] The foregoing are examples of ways in which the L2 controller **320** receives and processes various types of transactions from various types of masters, including non-coherent masters; coherent, non-caching masters; and coherent, caching masters. By handling such diverse combinations of transactions and master requirements in a single, unified controller, overall system flexibility is enhanced.

[0085] As explained, there is a need for the L2 cache subsystem **306** to include hardware, control logic, and/or state information to allow the L2 controller at **320** to accurately track and process the state of coherent, cache lines in the lower-level L1 cache subsystem **304**. In this example, the L1 cache subsystem **304** is utilizing a heterogeneous cache system, including the L1 main cache **314** and the L1 victim cache **316**. Examples of this disclosure allow the L2 controller **320** to maintain appropriate state information to accurately track the state of all coherent cache lines present in both the L1 main cache **314** and L1 victim cache **316**.

[0086] FIG. 5 shows an example of the L1 main cache **314** and the L1 victim cache **316**. In this example, as explained above, the L1 main cache **314** is a direct mapped cache, which thus has one way (Way **0**) and sets **0** through M. Continuing this example, as explained above, the L1 victim cache **316** is a fully associative cache, which thus has one set (Set **0**) and ways **0** through X.

[0087] FIG. 6 shows an example of the shadow L1 main cache **326** and the shadow L1 victim cache **328**, contained in the L2 cache subsystem **306**. The shadow L1 main cache **326** is a shadow copy of the address tag and MESI state information for the cache lines held in the L1 main cache **314**. The maintenance of this shadow copy enables the L2 controller **320** to track the lines that are cached in the L1 main cache **314**, for example to correctly decide when to send snoop transactions to either read or invalidate cache lines in the L1 main cache **314**. In this example, the shadow L1 main cache **326** also has one way (Way **0**) and sets **0** through M, permitting the shadow L1 main cache **326** to reflect the L1 main cache **314**.

[0088] The shadow L1 victim cache **328** is a shadow copy of the address tag and MESI state information for the cache lines held in the L1 victim cache **316**. As above with respect to the shadow L1 main cache **326**, the maintenance of the shadow L1 victim cache **328** enables the L2 controller **320** to accurately determine when to send snoop transactions to the L1 controller **310**. For example, if the shadow tags were not maintained in the L2 cache subsystem **306**, then the L2 controller **320** would need to snoop the L1 cache subsystem **304** for each request that could possibly be held in the L1 main or victim caches **314**, **316**, which could reduce performance due to the resulting snoop traffic bandwidth. In this example, the shadow L1 victim cache **328** includes one set (Set **0**) and ways **0** through X, along with floating entries, which render the shadow L1 victim cache **328** to reflect more entries than can be stored in the L1 victim cache **316**. The floating entries are explained in further detail below.

[0089] In both the shadow L1 main cache **326** and the shadow L1 victim cache **328**, only the tag (e.g., address) and coherence state information is shadowed. That is, in at least this example, it is not necessary to shadow the cached data itself.

[0090] When the L2 controller **320** receives a snoop transaction or a read or write transaction

occurs from the L3 controller **310** to the L2 controller **320**, the L2 controller **320** first checks the shadow L1 main and shadow L1 victim caches **326**, **328**. If a match is found (e.g., a hit), then the L2 controller **320** initiates a snoop transaction to the L1 controller **310**. When the snoop transaction returns, the L2 controller **320** uses the snoop response to update the shadow L1 main and shadow L1 victim caches **326**, **328**, if necessary.

[0091] Similarly, when the L1 controller **310** allocates a line in its L1 main cache **314**, or moves or relocates a line from the L1 main cache **314** to the L1 victim cache **316**, the L1 controller **310** communicates such movement to the L2 controller **320** to enable the L2 controller **320** to update the shadow L1 main and shadow L1 victim caches **326**, **328**. When the L1 controller **310** evicts a line from either the L1 main cache **314** or the L1 victim cache **316**, the line is either modified (e.g., dirty) or unmodified (e.g., clean) with respect to main memory (e.g., DDR **110**). The L1 controller **310** is configured to communicate both clean line evictions and dirty line victims to the L2 controller **320**, which enables the L2 controller **320** to accurately update its shadow L1 main and shadow L1 victim caches **326**, **328**. The signaling protocol to communicate such movement, relocation, and evictions between the L1 controller **310** and the L2 controller **320** is discussed in further detail below.

[0092] In an example, the L2 controller **320** learns that the L1 controller **310** is kicking a line out of its L1 victim cache **316** (e.g., to make room for a line coming from the L1 main cache **314**) before the L2 controller **320** receives the displaced victim from the L1 victim cache **316**. The line kicked out of the L1 victim cache **316** is held in a victim buffer **702** (e.g., as shown in FIG. 7) prior to being sent to the L2 controller **320** across the interface between the two controllers **310**, **320**. During this time period, the L2 controller **320** is aware of the transfer of a line from the L1 main cache **314** to the L1 victim cache **316**, which the L2 controller will cause to be mirrored in the shadow L1 main and shadow L1 victim caches **326**, **328**. However, the L2 controller **320** has not yet received the displaced victim from the L1 victim cache **316**, as the displaced victim is still in the victim buffer **702**.

[0093] The floating entries in the shadow L1 victim cache **328** address this issue. These floating entries extend the size of the shadow L1 victim cache **328** to include at least the number of victim buffers in the L1 cache subsystem **304**. In one example, the floating entries result in the shadow L1 victim cache **328** having twice the number of entries as the L1 victim cache **316**. In an example, the exact location of entries in the L1 victim cache **316** does not need to match the location of the same cache line as it is shadowed in the shadow L1 victim cache **328**. Decoupling the locations between the L1 victim cache **316** and the shadow L1 victim cache **328** improves the safety of the protocol, as a full address comparison is performed when the L2 controller **320** looks for an entry in the L1 victim cache **316**. Subsequently, when the L2 controller **320** receives the displaced victim across the interface from the victim buffer, the L2 controller **320** causes the line to be removed from its shadow L1 victim cache **328**.

[0094] FIG. 7 shows an example of an L1 cache subsystem **304** allocation of a new line at address C (e.g., line C), both before and after the allocation takes place. FIG. 8 shows the corresponding example from the view of the L2 cache subsystem **306**. Referring first to FIGS. 7 and 8 at once, before the allocation takes place, the L1 main cache **314** contains a cache line A that is in the modified (M) state, and the L1 victim cache **316** contains a cache line B that is also in the modified state. At the same time, the shadow L1 main cache **326** also contains the cache line A (e.g., tag and MESI data for the cache line A), which is in the same relative physical location within the shadow L1 main cache **326** as the cache line A in the L1 main cache **314**. Similarly, the shadow L1 victim cache **328** also contains the cache line B (e.g., tag and MESI data for the cache line B), which is not necessarily in the same relative physical location within the shadow L1 victim cache **328** as the cache line B in the L1 victim cache **316**.

[0095] When the L1 controller **310** decides to allocate line C, the L1 controller **310** conveys this allocation to the L2 controller (e.g., as part of a read request issued by the L1 controller **310**). In

this example, the address of line C maps to the same location in the L1 main cache **314** as the line A, and thus the L1 controller **310** relocates line A to the L1 victim cache **316**, in a location occupied by the line B. As a result of the line B being modified, the L1 controller **310** determines to send line B to the L2 cache subsystem **306** as a victim and moves the line B to the victim buffer **702**. After the read allocate for the line C, the L1 main cache **314** contains the line C in the location that formerly held the line A, the L1 victim cache **316** contains the cache line A that was relocated from the L1 main cache **314**, and the victim buffer **702** contains the cache line B that was evicted from the L1 victim cache **316**.

[0096] Similarly, after the read allocate for the line C (e.g., communicated by the L1 controller **310** to the L2 controller **320** as part of the read request for the line C), the shadow L1 main cache **326** contains the line C in the location that formerly held the line A and the shadow L1 victim cache **328** contains the relocated line A in one of its floating entries, while the line B also remains in the shadow L1 victim cache **328**. As explained above, there is a period of time in which the L2 controller **320** is aware that the L1 controller is moving the line A from the L1 main cache **314** to the L1 victim cache **316**, but the L2 controller **320** has not yet received the line B as a victim (e.g., the line B is still in the victim buffer **702**). The floating entries of the shadow L1 victim cache **328** provide an additional storage buffer, and the L2 controller **320** is configured to remove the line B from the shadow L1 victim cache **328** when the line B is received as a victim on the interface between the L2 cache subsystem **306** and the L1 cache subsystem **304**.

[0097] In general and as explained above, the L2 controller **320** is configured to receive an indication from the L1 controller **310** that a cache line is being relocated from the L1 main cache **314** to the L1 victim cache **316** (e.g., the cache line A in the example of FIGS. 7 and 8). In response to receiving the indication, the L2 controller **320** updates the shadow L1 main cache **326** to reflect that the cache line A is no longer located in the L1 main cache **314**. Similarly, in response to receiving the indication, the L2 controller **320** updates the shadow L1 victim cache **328** to reflect that the cache line A is located in the L1 victim cache **316**. The signaling protocol by which the L1 controller **310** communicates movement of cache lines between its L1 main cache **314**, L1 victim cache **316**, and victim buffer **702** are explained in further detail below. However, in one example the indication from the L1 controller **310** is a response to a snoop request from the L2 cache subsystem **306** to the L1 cache subsystem **304**. In another example, the indication from the L1 controller **310** is a read request from the L1 cache subsystem **304** to the L2 cache subsystem **306**.

[0098] These examples, in particular the floating entries of the shadow L1 victim cache **328**, enable cleaner handoff of a victim line from the L1 cache subsystem **304** to the L2 cache subsystem **306** by removing the timing window where a line is removed from the L1 victim cache **316**, but has not yet been received by the L2 cache subsystem **306** as a victim. Additionally, the L2 controller **320** maintaining accurate shadows of the L1 main cache **314** and the L1 victim cache **316** allows the L2 controller to only generate snoop transactions when necessary (e.g., when the L2 controller **320** is aware that a line is held in one of the L1 caches **314**, **316**).

[0099] As explained above, the L1 controller **310** communicates movement of cache lines between its L1 main cache **314**, L1 victim cache **316**, and victim buffer **702** to the L2 controller **320**. In some examples, this communication occurs in conjunction with a response to a snoop request from the L2 cache subsystem **306** to the L1 cache subsystem **304**. In other examples, this communication occurs in conjunction with a read request from the L1 cache subsystem **304** to the L2 cache subsystem **306**.

[0100] Referring back to FIG. 3, in some examples a transaction bus or interface between the L1 cache subsystem **304** and the L2 cache subsystem **306** contains a greater bandwidth than is needed to pass a transaction between the subsystems **304**, **306**. The transaction bus is represented schematically by the coupling between the L1 cache subsystem **304** and the L2 cache subsystem **306** (or similar couplings between L1 and L2 structures in FIGS. 1 and 2). The transaction bus has a bandwidth of $m+n$ bits, while a transaction (e.g., a read, a write, a snoop, a victim) only requires

m bits, leaving n bits of the transaction bus unused. Examples of this disclosure leverage this excess bandwidth on the transaction bus between the L1 cache subsystem **304** and the L2 cache subsystem **306** to communicate information from the L1 controller **310** to the L2 controller **320** in order to allow the L2 controller **320** to maintain its shadow L1 main cache **326** (e.g., tag and MESI information corresponding to the L1 main cache **314**) and shadow L1 victim cache **328** (e.g., tag and MESI information corresponding to the L1 victim cache **316**).

[0101] In particular, the L1 controller **310** is configured, in some examples, to send sideband signals in conjunction with a functional read transaction to the L2 controller **320**. The sideband signals contain information related to cache line movement (e.g., as described above with respect to the example of FIGS. **7** and **8**) occurring in the L1 cache subsystem **304**. Thus, the cache line movement information is communicated in parallel (e.g., as a part of a single transaction) with the functional read transaction that causes the cache line movement(s). The L2 controller **320** not only responds to transactions and information from the L1 controller **310**, but the L2 controller **320** also creates and enforces snoop transactions as required to maintain I/O (e.g., direct memory access (DMA)) coherence from non-caching requestors within the system (e.g., other CPU cores **102** in the system **100** may initiate a DMA request that is passed to the L2 controller **320** from a L3 controller, shared across CPU cores **102** as shown in FIG. **1**). In examples, these snoop transactions also cause the L2 controller **320** to initiate changes to its shadow L1 main cache **326** and shadow L1 victim cache **328**, as well as the L1 main cache **314** and the L1 victim cache **316**. For example, if the L1 controller **310** invalidates a line as a result of a snoop transaction (e.g., because the snoop transaction required invalidation, or because of a requirement due to the current state of the L1 main cache **314** or L1 victim cache **316**), the snoop response will indicate that the line transitioned to the invalid state. The L2 controller **320** then uses this information to update its shadow L1 main cache **326** or shadow L1 victim cache **328**. Thus, in addition to functional read transactions, the L1 controller **310** is configured to send additional sideband signals in conjunction with a response to a snoop transaction.

[0102] Examples of this disclosure reduce bandwidth on the transaction bus by avoiding the need for multiple messages to communicate both the functional read transaction and movements of cache lines within the L1 cache subsystem **304** that will result from that read transaction. Further, examples of this disclosure reduce timing dependencies and implementation complexity by avoiding the use of a separate asynchronous interface to communicate cache line movement information.

[0103] FIG. **9** shows a table 900 of sideband signaling protocol data in accordance with an example of this disclosure. The scope of this disclosure is not limited to any particular arrangement of signals within a transaction bus. For a given read transaction, the L1 controller **310** indicates to the L2 controller **320** whether the read transaction will allocate (the alloc signal) into the L1 main cache **314**, and if so, which line is moving from the L1 main cache **314** to the L1 victim cache **316**, and which line is moving out of the L1 victim cache **316**. If the alloc signal is de-asserted, then the L2 controller **320** disregards the remaining sideband signals.

[0104] In the table 900, the main_valid and victim_valid signals indicate whether the other main* and victim* signals, respectively, are valid. For example, the L1 controller **310** is configured to de-assert the valid signals when transmitted in parallel with a transaction that does not result in cache line movement(s) in the L1 main cache **314** and the L1 victim cache, respectively. The main_mesi and victim_mesi signals indicate the cache coherence state (e.g., MESI state) for a cache line moving from the L1 main cache **314** to the L1 victim cache **316** and for a cache line moving out of the L1 victim cache **316**, respectively. The main_secure and victim_secure signals indicate whether the cache line moving from the L1 main cache **314** to the L1 victim cache **316** and the cache line moving out of the L1 victim cache **316**, respectively, is secure. The main_address and victim_address signals indicate the addresses for the cache line moving from the L1 main cache **314** to the L1 victim cache **316** and for the cache line moving out of the L1 victim cache **316**,

respectively.

[0105] The L2 controller **320** is thus configured, in this example, to receive, in a single transaction, a read request in parallel with the aforementioned sideband signals that detail the cache line movement(s) that will occur in the L1 cache subsystem **304** as a result of the read request. In order for the L1 controller **310** to allocate space for data returned in response to the read request, the sideband signals indicate an address and coherence state of the cache line moving from the L1 main cache **314** to the L1 victim cache **316** and for the cache line moving out of the L1 victim cache **316**.

[0106] The L2 controller **320** is configured to update the shadow L1 main cache **326** to reflect that the cache line moving from the L1 main cache **314** to the L1 victim cache **316** is no longer present in the L1 main cache **314**. Similarly, the L2 controller **320** is configured to update the shadow L1 victim cache **328** to reflect that the cache line moving from the L1 main cache **314** to the L1 victim cache **316** is now present in the L1 victim cache **316**. If one or more of the valid bits in the sideband signals **900** are de-asserted, the L2 controller **320** is configured not to update its shadow L1 main cache **326** (main_valid de-asserted) or its shadow L1 victim cache **328** (victim_valid de-asserted).

[0107] In some examples, the L2 controller **320** is also configured to update the shadow L1 victim cache **328** to reflect that a cache line is no longer located in the L1 victim cache **316**. In particular, if the victim_mesi signal indicates that the cache line moving out of the L1 victim cache **316** has a coherence state other than modified (e.g., exclusive or shared), then the L2 controller **320** does not expect to receive a corresponding victim transaction because it is not necessary to write back a cache line that is not dirty. On the other hand, if the victim_mesi signal indicates that the cache line moving out of the L1 victim cache **316** has a modified coherence state, then the L2 controller **320** waits to receive a victim transaction (e.g., from the victim buffer **702**). Upon receiving the victim transaction, the L2 controller **320** is configured to update the shadow L1 victim cache **328** to reflect that a cache line is no longer located in the L1 victim cache **316**.

[0108] The foregoing examples reduce bandwidth on the transaction bus between the L1 cache subsystem **304** and the L2 cache subsystem **306** by avoiding the need for multiple messages to communicate both the functional read transaction and the movements of cache lines within the L1 cache subsystem **304** that will result from that read transaction.

[0109] The sideband signaling protocol discussed above leverages unused bandwidth on a transaction bus to facilitate communication of both the functional read transaction and the movements of cache lines within the L1 cache subsystem **304** that will result from that read transaction. However, in certain cases, the L1 controller **310** makes changes to the L1 main cache **314** and/or L1 victim cache **316** that are not coupled to a transaction that would be communicated to the L2 controller **320**. In these cases, the L2 controller **320** needs to be made aware of the changes to L1 main cache **314** and/or L1 victim cache **316** in another way.

[0110] In particular, for accurate coherent behavior, the L2 controller **320** maintains an accurate directory of the lines held in the L1 main cache **314** and L1 victim cache **316** (e.g., as shadow copies). This enables the L2 controller **320** to send snoop transactions to the L1 controller **320** to get the most up to date copy of the data when the L2 controller **320** knows the line is cached in the L1 cache subsystem **304**.

[0111] When the L1 controller **310** determines it must evict a non-modified line from the L1 victim cache **316** (e.g., for various reasons dependent on workload), the L1 controller **310** is configured in an example to inform the L2 controller **320** that the line is no longer held in the L1 cache subsystem **304**. In an example, the L1 controller **310** does not inform the L2 controller **320** that the line is no longer held in the L1 cache subsystem **304**. If the L1 controller **310** did not notify the L2 controller **320** that the line is no longer present, the L2 controller **320** may send at least one more snoop transaction to the address at a later time, believing that the line is still held in the L1 cache subsystem **304**. When the line is not found, the L1 controller will return a snoop response indicating that the line was not present. This concept is described as a snoop miss, and results in

unnecessary delays when the line was evicted knowingly by the L1 controller.

[0112] Examples of this disclosure address the foregoing by utilizing a tag update bus to employ a single cycle, pulsed protocol that enables the L1 controller **310** to communicate with the L2 controller **320** outside of the transaction-based sideband signaling protocol explained above. The tag update bus is separate from the transaction bus described above. Similarly to the transaction bus, the tag update bus is represented schematically by the coupling between the L1 cache subsystem **304** and the L2 cache subsystem **306** (or similar couplings between L1 and L2 structures in FIGS. **1** and **2**). Further, unlike transactions received over the transaction bus, which are held in a buffer and arbitrated before being utilized by the L2 controller **320**, the information provided over the tag update bus is usable by the L2 controller **320** upon receipt. The tag update bus protocol allows the L2 controller **320** to accurately maintain the shadow L1 main cache **326** and the shadow L1 victim cache **328**. In some cases, the tag update bus protocol is in the form of parallel signal groups, allowing the L1 controller **310** to communicate two or more cache updates to the L2 controller **320** per cycle.

[0113] By communicating the invalidations to the L2 controller **320**, unnecessary snoop transactions can be avoided, resulting in shorter latencies for processing transactions in the L2 cache subsystem **306**. Additionally, power savings may be realized by reducing the number of RAM accesses required by multiple arbitrations for the command that resulted in a snoop miss.

[0114] FIG. **10** shows a table 1000 of tag update bus protocol data in accordance with an example of this disclosure. The scope of this disclosure is not limited to any particular arrangement of signals within the tag update bus. In the table 1000, the t0_req and t1_req signals indicate whether the other t0 and t1 signals, respectively, are valid for use. When the L2 controller **320** detects that the t0_req or t1_req signals are asserted, the L2 controller **320** processes the remaining tag update bus signals. The t0_address and t1_address signals indicate the addresses for the cache line having its coherence state updated. The t0_mesi and t1_mesi signals indicate the cache coherence state (e.g., MESI state) for the cache line identified by t0_address and t1_address, respectively. The main_secure and victim_secure signals indicate whether the cache line identified by t0_address and t1_address, respectively, is secure.

[0115] In certain examples, t0_address and t1_address comprises an address in either the L1 main cache **314** or the L1 victim cache **316**, while in other examples the tag update bus is used solely to communicate updates to coherence state information for cache lines in the L1 victim cache **316**. In some examples, t0_mesi and t1_mesi could indicate any final cache coherence state for the cache line identified by t0_address and t1_address. The tag update bus provides the L1 controller **310** a means to communicate the cache line invalidations that result from the L1 controller **310**, while avoiding the snoop miss scenario described above.

[0116] The L2 controller **320** is thus configured to receive, over the tag update bus separate from a transaction bus, a message from the L1 controller **310** that includes a valid signal (e.g., t0_req), an address (e.g., t0_address), and a cache coherence state (e.g., t0_mesi). The message thus details an update to cache line coherence state(s) that will occur in the L1 cache subsystem **304**.

[0117] The L2 controller, in response to the valid signal being asserted, is configured to identify an entry in the shadow L1 main cache **326** or the shadow L1 victim cache **328** having an address corresponding to the address of the message and update a coherence state of the identified entry to be the coherence state of the message. In one example, the L2 controller **320** is configured only to identify an entry in the shadow L1 victim cache **328** having the address of the message.

Concurrently, the L2 controller **320** may receive transactions over the transaction bus from the L1 controller **310**. These transactions are separate from the message received over the tag update bus.

[0118] In some examples, the L2 cache subsystem **306** includes a transaction first-in, first-out buffer (FIFO, not shown for simplicity) coupled to the transaction bus that stores transactions received from the L1 cache subsystem **304** pending processing of those transactions by the L2 controller **320**. Messages received by the L2 controller **320** over the tag update bus are not stored in

the transaction FIFO, and are instead processed by the L2 controller **320** upon receipt of an asserted valid signal (e.g., `t0_req`).

[0119] In accordance with some examples of this disclosure, the L2 controller **320** exists in a system-wide MESI cache coherence implementation as described above. However, the L2 controller **320** is configured to remap certain transactions from certain masters to implement a local MEI protocol between the L2 controller **320** and the L1 controller **310** or the L3 controller **309**. In certain circumstances, remapping from MESI to MEI by the L2 controller **320** enables higher performance on read/write software workloads where memory locations are frequently read before being written to. For example, in a multi-core coherence solution, multiple caches can hold a cache line in the shared state. When one cache needs to modify the line, it must first initiate messaging to a downstream (e.g., numerically higher) controller which results in each of the other caches receiving an invalidating snoop to remove their copy. Finally, once snoop responses have been received, the initiating cache updates the cache coherence state of the line from shared to exclusive. The initiating cache then performs its cache line write and transitions the cache line to the modified state. Thus, when a line is actively shared between multiple different caches, and modified frequently, the number of coherence messages (read, write, victim, snoop) that are required can become large, negatively impacting the performance of software executing on the CPU core **302**. Suppression of the shared state by the L2 controller **320** causes each cache line allocation to bring the line in the exclusive state, reducing the need for future coherent messaging when a modification of the cache line occurs

[0120] FIG. **11** shows a block diagram of an exemplary flow **1100** of a transaction from the L1 controller, shown here as block **1102**, to the L2 instruction pipeline **1112**, prior to being processed by the L2 controller **320**. In the example of FIG. **11**, it is assumed that the transaction originates from the L1 controller **1102**; however, as will be explained further below, multiple masters could also issue the transaction. Regardless of the issuing master, the transaction is represented by block **1104** as a transaction that would invoke or generate a cache line in the shared coherence state.

[0121] In accordance with examples of this disclosure, the L2 controller **320** suppresses the shared state by applying opcode mapping logic **1106** to the transaction **1104**. As will be explained further below, the opcode mapping logic **1106** maps a transaction opcode to a subset of opcodes for the final coherent cache state of the cache line comprising the modified, exclusive, or invalid states. In particular, opcodes that would have resulted in a final cache line coherence state of shared are remapped to one of this subset of opcodes. The opcode mapping logic **1106** need not map opcodes that would have resulted in a final cache state of modified, exclusive, or invalid.

[0122] The first request, or transaction **1104**, is thus mapped to a second request demonstrated by block **1108**, which avoids invoking the shared coherence state. The second request is then arbitrated as normal by L2 arbitration logic **1110** and enters the L2 instruction pipeline **1112**, to be subsequently processed by the L2 controller **320**.

[0123] In some examples, the L2 cache subsystem **306** includes a configuration register shown as block **1107**, which includes a shared field. The shared field allows the L2 cache subsystem **306** to be programmatically configured to either suppress the shared state, or not to suppress the shared state (e.g., not perform opcode mapping and function as a part of the larger MESI-based coherence system, described above). For example, if the shared field in configuration register **1107** is asserted, then the opcode mapping logic **1106** is not employed to map transaction opcodes to that would have resulted in a final cache line coherence state of shared. Thus, if a transaction **1104** is received as a third request when the shared field is asserted, the third request is processed by the L2 controller **320** without having its opcode mapped by the opcode mapping logic **1106**.

[0124] FIG. **12** shows a flow chart of a method **1200** in accordance with various embodiments. The method **1200** begins in block **1202** with the L2 controller **320** receiving a first request for a cache line in a shared cache coherence state. The request may be received from the L1 controller **310** as a read request, from the streaming engine **210** as a CMO that requires the L2 controller **320** to issue a

snoop to the L1 controller **310**, or from the L3 controller **309** as a snoop that requires the L2 controller **320** to issue a snoop to the L1 controller **310**.

[0125] The method **1200** continues in block **1204** with the L2 controller **320** mapping the first request to a second request for a cache line in an exclusive cache coherence state, as explained above. For example, opcode mapping logic **1106** is applied to the opcode of the first request (e.g., invoking the shared coherence state) to map to the opcode of the second request (e.g., invoking the exclusive coherence state). As explained, the opcode mapping logic **1106** is carried out prior to the request entering the L2 arbitration logic **1110** and the L2 instruction pipeline **1112**, and thus being processed by the L2 controller **320**. In an example, read requests are either allocating or non-allocating, and either non-coherent or coherent. In this example, the opcode mapping logic **1106** maps non-coherent allocating reads to a read command without snoop, because no snooping is required for a non-coherent read and reading without snooping does not invoke the shared state. Similarly, the opcode mapping logic **1106** also maps non-coherent non-allocating reads to a read command without snoop. The opcode mapping logic **1106** maps coherent allocating reads to a read exclusive command, which guarantees that the line will be allocated in the exclusive state. The opcode mapping logic **1106** maps coherent non-allocating reads to a read once command, since these only need to sample the coherent data (e.g., not allocate), and thus the current owner can keep the line without invoking the shared state. In another example, certain snoop commands (e.g., from the L3 controller **309**) or CMOs have an opcode that would normally require a line to transition to the shared state. In this example, the opcode mapping logic **1106** maps such snoop commands and CMOs to a snoop command or CMO, respectively, that requires the line to instead transition to the invalid state. Additionally, if the L2 controller **320** determines to send a snoop command to the L1 controller **310**, the opcode mapping logic **1106** maps such a snoop command to a snoop command that requires the L1 controller **310** to instead transition the line to the invalid state.

[0126] The method **1200** then continues in block **1206** with the L2 controller **320** responding to the second request, if the second request is of a type that warrants a response (e.g., if the second request is a read response, a read response is warranted). In the event that the requested cache line is stored in the L2 cache subsystem **306**, as part of its response, the L2 controller **320** transitions a coherence state of the cache line to invalid rather than shared. Alternately, the method **1200** continues in block **1208** with forwarding the second request. For example, where the first request results in a snoop being issued by the L2 controller **320**, the L2 controller forwards the second request for the cache line in the exclusive state, rather than the shared state.

[0127] In some examples, the L1 controller **310** determines to change a size of the L1 main cache **314**. For example, the L1 main cache **314** may be an allocated region of the larger L1 SRAM **312** that can grow (e.g., from 32 KB to 64 KB) or shrink (e.g., from 32 KB to 16 KB) in size over time, depending on requirements communicated to the L1 controller **310**, for example from the CPU core **302** or software executing thereon. The L2 controller **320** needs to be aware of changes in size to the L1 main cache **314**, so that the L2 controller **320** can properly maintain (e.g., change the size of) its shadow L1 main cache **326**.

[0128] The following protocol enables the L2 controller **320**, in an example, to change the size of its shadow L1 main cache **326** while avoiding data corruption and/or transaction deadlocks (e.g., where a first transaction relies on a second transaction, which is pending resolution of the first transaction). In one example, sideband signals of the transaction bus (explained above) are used by the L1 controller **310** to communicate the size change of the L1 main cache **314** to the L2 controller **320**. In this example, reference is made to certain ones of the sideband signals of the transaction bus, in particular referred to as: `global_on`, `global_coh_type`, and `cache_size`. The `global_on` signal indicates that the L1 controller **310** is performing a global operation on its L1 main cache **314**. The `global_coh_type` signal indicates the type of global coherence operation being performed on the L1 main cache **314**. In the specific example of a size change of the L1 main cache **314**, the `global_coh_type` will be a writeback invalidate operation. During a cache size change, coherence is

maintained by writing the data to the endpoint and by invalidating the cacheline. The `cache_size` signal indicates the size to which the L1 main cache **314** is transitioning.

[0129] FIG. **13** shows a flow chart of a method **1300** for changing the size of the L1 main cache **314**, and the resultant change in size of the shadow L1 main cache **326**. The method **1300** begins in block **1302** with determining, by the L1 controller **310**, to change a size of the L1 main cache **314**. This determination to change the cache size is, for example, the result of a control or configuration register write programming a configuration register of the L1 controller **310** to indicate the desired new cache size and initiate the cache size change.

[0130] The method **1300** continues in block **1304** with the L1 controller **310** servicing pending read and write requests from a CPU core, such as the CPU core **302**. The method **1300** then continues in block **1306** in which the L1 controller **310** stalls new read and write requests from the CPU core **302**. This allows the L1 controller **310** to work through pending requests but restrict new requests so that it may perform the global operation (e.g., writeback invalidate and cache size change) on the L1 main cache **314**.

[0131] The method **1300** continues in block **1308** with the L1 controller **310** writing back and invalidating the L1 main cache **314**. At this point in the method **1300**, the L1 controller **310** asserts the `global_on` signal to indicate it is performing a global operation, and the `global_coh_type` indicates a writeback invalidate as explained above. The L1 controller **310** is configured to send victims to the L2 controller **320** during this stage, which enables the L2 controller **320** to update the shadow L1 main and victim caches **326**, **328**. If the victim hits in the L2 cache **324**, the L2 controller **320** is also configured to update that cache line with the victim data. If the victim is not present in L2 cache **324**, the L2 controller **320** forwards the victim (e.g., to the L3 controller **309**). During the size change of the L1 main cache **314**, coherence is maintained writing the data back to the endpoint and invalidating the cache line. While the L1 controller **310** writes back and invalidates the L1 main cache **314**, the L1 controller **310** is also configured to accept and stall a snoop request from the L2 controller **320**.

[0132] While the L1 controller **310** asserts the `global_on` signal (e.g., during a global operation), the L1 controller **310** also de-asserts a ready signal, which indicates to the CPU core **302** not to send the L1 controller **310** additional requests for a cache size change or other global coherence operations. The ready signal remains de-asserted until the global operation is completed (e.g., the `global_on` signal is de-asserted).

[0133] Once the `global_on` signal is de-asserted, the L1 controller **310** responds to any pending snoop transactions that were received from the L2 controller **320** and stalled by the L1 controller **310** during the writeback invalidate (e.g., the global coherence operation for L1 main cache **314** size change). In an example, the L1 controller **310** responds to the pending snoop transactions with a response indicating a cache miss because the L1 main cache **314** is invalidated as part of the size change protocol. Once the `global_on` signal is de-asserted, the L1 controller **310** also begins accepting read and write requests from the CPU core **302** using the new cache size for the L1 main cache **314**. At this point the L1 controller **310** has implemented the functionality to change the size of its L1 main cache **314**.

[0134] The method **1300** then continues to block **1310** in which the L2 controller **320** receives an indication that the L1 main cache **314** has been invalidated and had its size changed. In an example, the L1 controller **310** sends such an indication to the L2 controller in response to the L1 controller **310** having received write responses for all victims written back by the L1 controller **310**, while no further victims are pending to be written back by the L1 controller **310**.

[0135] In this example, the L1 controller **310** uses sideband signals of `global_on`, `global_coh_type`, and `cache_size` to communicate that the L1 main cache **314** has been invalidated and had its size changed. For example, when `global_coh_type` indicates a writeback invalidate and the `cache_size` signal has changed, the L1 controller **310** de-asserting `global_on` indicates to the L2 controller **320** that the L1 main cache **314** has been invalidated and had its size changed. This indication allows

the L2 controller **320** to begin the process of resizing its shadow L1 main cache **326**. To begin resizing the shadow L1 main cache **326**, the L2 controller **320** flushes its pipeline, or completes all transactions that are present in its pipeline while stalling transactions from other masters. [0136] In some examples, the L2 controller **320** flushes its pipeline in separate phases, which include a blocking soft stall phase, a non-blocking soft stall phase, and a hard stall phase. In general, blocking transactions include read requests and write requests that are not victims, which have the potential to create a secondary transaction (e.g., a snoop), while non-blocking transactions include victims, snoops, and all responses.

[0137] In an example, during the blocking soft stall phase, the L2 controller **320** stalls all blocking transactions, such as fetches, read requests, and write requests from the CPU core **302** and DMA read/write accesses (e.g., from another CPU core) but allows response transactions, non-blocking snoop and victim transactions to be accepted and arbitrated. In some examples the L2 controller **320** flushes its pipeline over several cycles. Following the blocking soft stall phase, the L2 controller **320** enters the non-blocking soft stall phase, in which the L2 controller **320** allows response transactions and victims but stalls snoop transactions, in addition to the blocking transactions already stalled in the blocking soft stall phase. As a result, the L2 controller **320** does not initiate new snoops to the L1 main cache **314** for lines previously cached in in the L1 main cache **314**.

[0138] After the L2 controller **320** pipeline is flushed, the method **1300** continues to block **1312** in which the L2 controller **320** stalls requests received from any master. This phase is the hard stall phase referred to above. In particular, the L2 controller **320** pipeline is flushed, the L2 controller **320** enforce a hard stall where all transactions, including response transactions, are stalled from all masters.

[0139] In some examples, the L2 controller **320** also de-asserts or causes the ready signal (explained above with respect to the L1 controller) to be de-asserted. By de-asserting the ready signal, the L2 controller **320** prevents the CPU core **302** from sending requests for a cache size change or other global coherence operation until the L2 controller **320** has completed the currently-pending request (e.g., a cache size change). In one example, the ready signal provided to the CPU core **302** comprises a logical AND of a ready signal from the L1 controller **310** and the L2 controller **320**. That is, the CPU core **302** only receives an asserted ready signal when both the L1 controller **310** and the L2 controller **320** assert their ready signals (e.g., when the cache size change operation is complete).

[0140] When the hard stall is enforced in block **1312**, the method **1300** then continues to block **1314** in which the L2 controller **320** reinitializes the shadow L1 main cache **326** to clear its previous contents (e.g., invalidate cache lines previously held in the shadow L1 main cache **326**) and change a size of the shadow L1 main cache **326**. In some examples, reinitializing the shadow L1 main cache **326** takes several cycles, during which the L2 controller **320** continues to enforce the hard stall on other masters. Once the shadow L1 main cache **326** is reinitialized, the L2 controller **320** unstalls the masters and asserts its ready signal. The L2 controller **320** then begins to process pending transactions from one or more holding buffers, and accepts new transactions. At this point the size change protocol execution is complete. In some cases, the L1 controller **310** sends a transaction (e.g., a read request) to the L2 controller **320** while the L2 controller **320** is flushing its pipeline in block **1310** or stalled in block **1312**, and thus the transaction from the L1 controller **310** is stalled as well. The L2 controller **320** responds to such transactions after reinitializing the shadow L1 main cache **326**.

[0141] In the foregoing discussion and in the claims, the terms “including” and “comprising” are used in an open-ended fashion, and thus should be interpreted to mean “including, but not limited to . . .” Also, the term “couple” or “couples” is intended to mean either an indirect or direct connection. Thus, if a first device couples to a second device, that connection may be through a direct connection or through an indirect connection via other devices and connections. Similarly, a

device that is coupled between a first component or location and a second component or location may be through a direct connection or through an indirect connection via other devices and connections. An element or feature that is “configured to” perform a task or function may be configured (e.g., programmed or structurally designed) at a time of manufacturing by a manufacturer to perform the function and/or may be configurable (or re-configurable) by a user after manufacturing to perform the function and/or other additional or alternative functions. The configuring may be through firmware and/or software programming of the device, through a construction and/or layout of hardware components and interconnections of the device, or a combination thereof. Additionally, uses of the phrases “ground” or similar in the foregoing discussion are intended to include a chassis ground, an Earth ground, a floating ground, a virtual ground, a digital ground, a common ground, and/or any other form of ground connection applicable to, or suitable for, the teachings of the present disclosure. Unless otherwise stated, “about,” “approximately,” or “substantially” preceding a value means ± 10 percent of the stated value. [0142] The above discussion is meant to be illustrative of the principles and various embodiments of the present disclosure. Numerous variations and modifications will become apparent to those skilled in the art once the above disclosure is fully appreciated. It is intended that the following claims be interpreted to embrace all such variations and modifications.

Claims

1. A method, comprising: receiving, by a cache controller of a first cache subsystem of a memory hierarchy, a first request to read first data, wherein the memory hierarchy includes the first cache subsystem and a second cache subsystem, wherein the first cache subsystem includes a first shadow cache corresponding to a first cache of the second cache subsystem; determining, by the cache controller, whether the first request is a non-coherent request or a coherent request; and based on determining that the first request is a coherent request, determining, by the cache controller, whether the first request is a hit or a miss in the first shadow cache; and based on determining that the first request is a hit in the first shadow cache, providing, by the cache controller, a snoop request for the first data from the first cache; and providing, by the cache controller, a response to the first request that includes the first data received from the first cache.
2. The method of claim 1, wherein the first cache is a level one (L1) main cache, and the first shadow cache is a shadow L1 main cache.
3. The method of claim 2, wherein the first cache is a L1 program cache.
4. The method of claim 3, wherein determining whether the first request is a non-coherent request or a coherent request comprises determining whether or not the L1 program cache is an initiator of the first request.
5. The method of claim 3, wherein the memory hierarchy further includes a level three (L3) cache, and wherein determining whether the first request is a non-coherent request or a coherent request comprises determining whether or not the L3 cache is an initiator of the first request.
6. The method of claim 2, wherein the first cache subsystem further comprises a second shadow cache corresponding to a second cache of the second cache subsystem, and wherein the method further comprises: based on determining that the first request is a coherent request, determining whether the first request is a hit or a miss in the second shadow cache; and based on determining that the first request is a hit in the second shadow cache, providing a snoop request for the first data in the second cache.
7. The method of claim 6, wherein the second cache is a level one (L1) victim cache, and the second shadow cache is a shadow L1 victim cache.
8. The method of claim 2, wherein the first cache subsystem further comprises a second cache, wherein the method further comprises: based on determining that the first request is a miss in the first shadow cache, retrieving the first data from the second cache.

9. The method of claim 8, wherein the second cache is a level two (L2) cache.

10. The method of claim 1, wherein determining whether the first request is a non-coherent request or a coherent request comprises determining whether the first request is a non-coherent request or a coherent request based on an initiator of first the request.

11. The method of claim 1, further comprising: receiving, by the cache controller, a second request to write second data; determining, by the cache controller, whether the second request is a hit or a miss in the first shadow cache; and based on determining that the second request is a hit in the first shadow cache, providing a snoop read and invalidate request to the first cache.

12. The method of claim 11, further comprising: receiving third data from the first cache in response to the snoop read and invalidate request; determining whether the third data is indicated as modified; and based on determining that the third data is indicated as modified, merging the second data and the third data to produce merged data; and causing the merged data to be written for the second request.

13. The method of claim 12, further comprising: based on determining that the third data is not indicated as modified, causing the second data to be written for the second request.

14. The method of claim 1, further comprising based on a response to the snoop request from the first cache not including valid data, retrieving the first data from an endpoint other than the first cache.

15. A system, comprising: a first cache subsystem in a first level of a memory hierarchy, the first cache subsystem comprising a first cache; a second cache subsystem in a second level of the memory hierarchy, the second cache subsystem comprising: a first shadow cache corresponding to the first cache of the first cache subsystem; and a cache controller capable of: receiving a first request to read first data; determining whether the first request is a non-coherent request or a coherent request; and based on determining that the first request is a coherent request, determining whether the first request is a hit or a miss in the first shadow cache; and based on determining that the first request is a hit in the first shadow cache, providing a snoop request for the first data from the first cache; and providing a response to the first request that includes the first data received from the first cache.

16. The system of claim 15, wherein the first level is level one (L1) and the first cache is a level one (L1) main cache, and wherein the second level is level two (L2) and the first shadow cache is a shadow L1 main cache.

17. The system of claim 16, wherein the first cache subsystem further comprises a second cache, wherein the second cache subsystem further comprises a second shadow cache, and wherein the cache controller is capable of: based on determining that the first request is a coherent request, determining whether the first request is a hit or a miss in the second shadow cache; and based on determining that the first request is a hit in the second shadow cache, providing a snoop request for the first data in the second cache.

18. The system of claim 17, wherein the second cache is a level one (L1) victim cache, and the second shadow cache is a shadow L1 victim cache.

19. The system of claim 15, wherein the second cache subsystem further comprises a second cache, and wherein the cache controller is capable of: based on determining that the first request is a miss in the first shadow cache, retrieving the first data from the second cache.

20. The system of claim 19, wherein the second cache is a L2 cache.
