

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265239

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Wang; Huai-Mu et al.

Quantization Error Compensation for Vector Computing

Abstract

A method for performing a computing task includes: extracting one or more features from a user content; converting the features to a floating point query vector; quantizing the floating point query vector; obtaining a database vector including one or more floating point feature vectors; determining a compensation vector based on a data distribution of the floating point query vector; quantizing the floating point feature vectors; determining an error function based on a difference between data distributions of i) the quantized query vector compensated with the compensation vector, and ii) the floating point query vector; determining, based on the error function, values of the compensation vector corresponding to the quantized feature vectors; combining the quantized query vectors and the values of the compensation vector to obtain one or more compensated query vectors; and performing the computing task using the compensated query vectors and the quantized feature vectors to obtain an output.

Inventors: Wang; Huai-Mu (New Taipei, TW), Hu; Han-Wen (Zhubei, TW), Lin; Bo-Rong (Taichung, TW)

Applicant: Macronix International Co., Ltd. (Hsinchu, TW)

Family ID: 1000007774282

Assignee: Macronix International Co., Ltd. (Hsinchu, TW)

Appl. No.: 18/583216

Filed: February 21, 2024

Publication Classification

Int. Cl.: G06F16/23 (20190101); G06F16/22 (20190101); G06F16/9535 (20190101)

U.S. Cl.:

CPC G06F16/2365 (20190101); G06F16/2237 (20190101); G06F16/9535 (20190101);

Background/Summary

TECHNICAL FIELD

[0001] The present disclosure is related to quantization error compensation for vector computing, e.g., in-memory computing (IMC) in semiconductor memory devices.

BACKGROUND

[0002] Vector computing technology is used in applications such as artificial intelligence (AI) and machine learning (ML) applications. Vector computing can take place in processors or partially or completely in a memory device. Computing within a memory device is referred to as IMC. Compared to traditional data processing techniques that access data from memory devices and perform computing in a separate circuit, IMC allows the computing to take place within the memory devices, which can increase data processing speed.

SUMMARY

[0003] The present disclosure provides techniques relating to quantization error compensation for vector computing. Some semiconductor devices in which the disclosed techniques can be applied include non-volatile memory devices, such as NAND or NOR flash memory devices.

[0004] In one aspect, the disclosure describes a method for performing a computing task based on a user content. The method includes extracting one or more features from the user content, the one or more features corresponding to a search request indicated in the user content. The method includes converting the one or more features to a floating point query vector. The method includes quantizing the floating point query vector to obtain a quantized query vector. The method includes obtaining a database vector based on the one or more features, with the database vector including one or more floating point feature vectors. The method includes determining a compensation vector based on a data distribution of the floating point query vector. The method includes quantizing the one or more floating point feature vectors to obtain one or more quantized feature vectors. The method includes determining an error function based on a difference between data distributions of i) the quantized query vector compensated with the compensation vector, and ii) the floating point query vector. The method includes determining, based on the error function, one or more values of the compensation vector corresponding to the one or more quantized feature vectors. The method includes combining the one or more quantized query vectors and the one or more values of the compensation vector to obtain one or more compensated query vectors. The method includes performing the computing task using the one or more compensated query vectors and the one or more quantized feature vectors to obtain an output.

[0005] In some implementations, determining the one or more values of the compensation vector comprises: determining the one or more values of the compensation vector such that a difference in magnitude between the quantized query vector compensated with the compensation vector and each of the one or more quantized feature vectors, is below a known threshold value.

[0006] In some implementations, the computing task comprises performing at least one of a multiply-and-accumulate operation, general matrix multiplication (GeMM), fully connected layer computing, or k-nearest neighbors computing.

[0007] In some implementations, the computing task comprises computing a plurality of vector distances.

[0008] In some implementations, the plurality of vector distances comprises at least one of: a plurality of cosine similarity distances, a plurality of Euclidean distances, or a plurality of Hamming distances.

[0009] In some implementations, the method includes generating a response to the user content based on the output, wherein generating the response comprises sorting the plurality of vector distances in a descending order or an ascending order.

[0010] In some implementations, the user content comprises at least one of: graphical information, textual information, geographical information, or temporal information.

[0011] In some implementations, the method includes receiving the user content from at least one of: a text-based search engine, a graph-based search engine, a brute force search engine, or a behavior-based content recommendation system.

[0012] In some implementations, performing the computing task comprises performing the computing task using an in-memory computing (IMC) circuit, and wherein the IMC circuit comprises a plurality of memory cells comprising at least one of NAND flash cells, NOR flash cells, phase change memory (PCM), Magnetoresistive random-access memory (MRAM), Ferroelectric random-access memory (FeRAM), or Spin-Transfer-Torque random-access memory (STT-RAM).

[0013] In one aspect, the disclosure describes a computing system having one or more processors and a computing circuit coupled to the one or more processors. The one or more processors are configured to execute instructions to perform operations of the method described above.

[0014] In one aspect, the disclosure describes a non-transitory computer-readable medium storing program instructions that, when executed, cause one or more processors to perform operations of the method described above.

[0015] The details of one or more disclosed implementations are set forth in the accompanying drawings and the description below. Other features, aspects, and advantages will become apparent from the description, the drawings and the claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0016] FIG. 1 illustrates a block diagram of an example system including a memory device, according to one or more implementations.

[0017] FIG. 2 illustrates a block diagram of an example memory device, according to some implementations.

[0018] FIG. 3 illustrates an example of quantization error compensation for a multiply-and-accumulate operation, according to some implementations.

[0019] FIG. 4 illustrates pseudo code of an example algorithm for quantization error compensation, according to some implementations.

[0020] FIG. 5 provides graphical representations comparing quantization error with and without compensation, according to some implementations.

[0021] FIGS. 6A-6B illustrate a flowchart of an example method, according to some implementations.

[0022] Figures are not drawn to scale. Like reference numbers refer to like components.

DETAILED DESCRIPTION

[0023] The rise of AI and ML applications is accompanied by increasing needs for large-scale, complex computing capacity. Some AI and ML applications model input data as vectors and use processors and/or memory devices to perform logic or arithmetic computation based on the vectors. For example, a host device, which serves as an edge AI device, can provide vectorized data to a memory device and instruct the memory device to perform certain computations. The memory

controller can then provide the vectorized data to the memory cells for computing and return the results back to the host device. By localizing computing tasks within the memory cells, the time for a processor (e.g., the memory controller) to return results can be largely reduced, which can lead to increase in computing efficiency and capacity.

[0024] In a more specific example, a host device can execute an AI application to perform a query-based search, e.g., search for a movie. The query can include textual, audio, video, or other types of content obtained from a user of the AI application. The AI application can use numeric formatted vectors to represent features of the query, such as the length, language, and genre of the movie. Upon obtaining the vectors from the AI application, a processor can either use its own computing circuitry or configure the memory cells to perform the computing, such as finding a closest match to the query.

[0025] Due to circuit restrictions, e.g., circuit bus bandwidth in terms of bits, floating point data provided by some applications are converted to integer type before being processed. For example, quantization may be performed to convert the floating point data to integer-typed data before processing, and dequantization may be performed to convert the processed data back to floating point. In converting from floating point to integer data, some precision can be lost due to the limitations of the number of bits available to represent digital integer data. Such loss in precision can cause quantization error, due to which the computing result can deviate from a result that would have been obtained if the computation was performed directly using floating point data without quantization. While quantization error can be reduced by introducing more bits to the quantization process to increase precision, the increase of precision can incur more circuit complexity and cost.

[0026] Implementations of this disclosure provide techniques that can reduce quantization error without requiring increased precision (e.g., higher number of bits) to represent integer data. As described in detail below, one or more implementations allow for obtaining a compensation vector to be combined with the input data vector after quantization. By compensating the input data vector, the computing circuitry can return results that are closer to the results that would have been returned should the computation be performed using floating point data. Using the results with reduced error, an application can provide the user with more accurate query results, such as a movie that more closely match the user's interest, which can improve the user's experience with the application.

[0027] The techniques can be applied to any suitable circuits or devices such as semiconductor devices. For illustration purposes, a memory device is described herein as an example of a semiconductor device. For example, the techniques can be applied to various types of non-volatile memory devices such as NOR flash memory devices, NAND flash memory devices, phase change memory (PCM), erasable programmable read-only memory (EPROM), Static Random Access Memory (SRAM), Resistive Random Access Memory (RRAM), Ferroelectric Random Access Memory (FeRAM), Magnetoresistive Random Access Memory (MRAM), Spin-Transfer-Torque Random Access Memory (STT-RAM) among others. The techniques can be applied to various memory types, such as SLC (single-level cell) devices, MLC (multi-level cell) devices like 2-level cell devices, TLC (triple-level cell) devices, QLC (quad-level cell) devices, or PLC (penta-level cell) devices. Additionally or alternatively, the techniques can be applied to various types of devices and systems, such as secure digital (SD) cards, embedded multimedia cards (eMMC), solid-state drives (SSDs), embedded systems, or computing network devices, among others.

[0028] FIG. 1 illustrates an example system **100** including a memory device, according to one or more implementations. The system **100** includes a memory device **110** and a host device **120**. The memory device **110** includes a device controller **112** and a memory **116**. The device controller **112** includes a processor **113** and an internal memory **114**. In some implementations, the memory device **110** includes a plurality of memories **116** that are coupled to the device controller **112**.

[0029] In some implementations, the memory device **110** is a storage device. For example, the

memory device **110** can be an embedded multimedia card (eMMC), a secure digital (SD) card, a solid-state drive (SSD), or some other suitable storage. In some implementations, the memory device **110** is a client device that is coupled to a host device **120**. For example, the memory device **110** is an SD card in a digital camera or a media player that is the host device **120**.

[0030] The device controller **112** is a general-purpose microprocessor, or an application-specific microcontroller. In some implementations, the device controller **112** is a memory controller for the memory device **110**. The following sections describe the various techniques based on implementations in which the device controller **112** is a memory controller. However, the techniques described in the following sections are also applicable in implementations in which the device controller **112** is another type of controller that is different from a memory controller.

[0031] The processor **113** is configured to execute instructions and process data. The instructions include firmware instructions and/or other program instructions that are stored as firmware code and/or other program code, respectively, in the secondary memory. The data includes program data corresponding to the firmware and/or other programs executed by the processor, among other suitable data. In some implementations, the processor **113** is a general-purpose microprocessor, or an application-specific microcontroller. The processor **113** is also referred to as a central processing unit (CPU).

[0032] The processor **113** accesses instructions and data from the internal memory **114**. In some implementations, the internal memory **114** is a Static Random Access Memory (SRAM) or a Dynamic Random Access Memory (DRAM). For example, in some implementations, when the memory device **110** is an eMMC, an SD card or a smart watch, the internal memory **114** is an SRAM. In some implementations, when the memory device **110** is a digital camera or a media player, the internal memory **114** is DRAM.

[0033] In some implementations, the internal memory is a cache memory that is included in the device controller **112**, as shown in FIG. 1A. The internal memory **114** stores instruction codes, which correspond to the instructions executed by the processor **113**, and/or the data that are requested by the processor **113** during runtime.

[0034] The device controller **112** transfers the instruction code and/or the data from the memory **116** to the internal memory **114**. The memory **116** can be a semiconductor device. In some implementations, the memory **116** is a non-volatile memory that is configured for long-term storage of instructions and/or data, e.g., a NAND flash memory device, or some other suitable non-volatile memory device. In implementations where the memory **116** is NAND flash memory, the memory device **110** is a flash memory device, e.g., a flash memory card, and the device controller **112** is a NAND flash controller. For example, in some implementations, when the memory device **110** is an eMMC or an SD card, the memory **116** is a NAND flash; in some implementations, when the memory device **110** is a digital camera, the memory **116** is an SD card; and in some implementations, when the memory device **110** is a media player, the memory **116** is a hard disk.

[0035] The memory **116** includes a plurality of blocks. The memory **116** can be a two-dimensional (2D) memory including 2D memory blocks. The memory **116** can also be a three-dimensional (3D) memory including 3D memory blocks.

[0036] The system **100** can be used for AI applications. For example, the host device **120** can be an edge AI device configured to execute an AI application that involves a computing task. The host device **120** can transmit data used in the computing task to the memory device **110**, which is configured to perform IMC in the memory **116**. The host device **120** and the memory device **110** can be integrated on the same piece of hardware or distributed across separate pieces of hardware.

[0037] FIG. 2 illustrates an example memory device **200**, according to one or more implementations. The memory device **200** can be implemented as the memory **116** of FIG. 1. The memory device **200** includes a memory cell array **210**. The memory cell array **210** can include a number of memory cells coupled in series to a number of row word lines and a number of column bit lines.

[0038] A memory cell can include a memory transistor configured as a storage element. The memory transistor can include a silicon-oxide-nitride-oxide-silicon (SONOS) transistor, a floating gate transistor, a nitride read only memory (NROM) transistor, or any suitable non-volatile memory metal-oxide-semiconductor (MOS) device that can store charges.

[0039] The memory device **200** includes a memory interface **202** having multiple input/output (I/O) ports for receiving data, e.g., from a controller such as the device controller **112** of FIG. **1** or the host controller **122** of FIG. **1**, or outputting data from the memory cell array **210**. The memory device **200** includes a data buffer **240** configured to buffer data through the memory interface **202**. The data buffer **240** can include a data input buffer **242** configured to buffer/transmit data from a controller (e.g., the device controller **112** of FIG. **1** or the host controller **122** of FIG. **1**) through the memory interface **202** to the memory cell array **210**. The data buffer **240** can also include a data output buffer **244** configured to buffer/transmit out data from the memory cell array **210** through the memory interface **202**, e.g., to a host device such as the host device **120** of FIG. **1**.

[0040] In some implementations, the memory device **200** further includes an X-decoder (or row decoder) **206** and an optional Y-decoder (not shown). Each memory cell is coupled to the X-decoder **206** via a respective word line and coupled to the Y-decoder via a respective bit line **215**. Accordingly, each memory cell can be selected by the X-decoder **206** and the Y-decoder for read or write operations through the respective word line and the respective bit line **215**.

[0041] The memory device **200** includes a page buffer circuit **220** that includes a number of page buffers. Each page buffer is connected to the memory cell array **210** through a respective bit line **215**. In some implementations, a page buffer is connected to the Y-decoder through a data line associated with a corresponding bit line **215** that connects a corresponding line of memory cells in the memory cell array **210**. A page buffer is configured to control a voltage on a corresponding bit line to perform an operation, e.g., read, program, or erase, on a memory cell coupled to the corresponding bit line. A page buffer can include at least one latch circuit.

[0042] In some implementations, the memory device **200** further includes a data cache circuit **230** coupled between the page buffer circuit **220** and the data buffer **240**. During a program or erase operation, the data cache circuit **230** is configured to store data from the data buffer **240** (e.g., from the data input buffer **242**) and/or output through the page buffer circuit **220** to the memory cell array **210**. During a read operation, the data cache circuit **230** is configured to store data from the memory cell array through the page buffer circuit **220** and/or output data to the data buffer **240** (e.g., to the data output buffer **244**).

[0043] In some implementations, the memory device **200** further includes a control logic **204** coupled to components in the memory device **200** including the X-decoder **206** and the Y-decoder, the data buffer **240**, the page buffer circuit **220**, and the data cache circuit **230**. The control logic **204** is configured to receive a command, address information, and/or data, e.g., from a memory controller such as the device controller **112** or the host controller **122** of FIG. **1**, via the memory interface **202**. The control logic **204** can also process the command, the address information, and/or the data, for example, to generate physical address information, e.g., of blocks/pages, in the memory cell array **210**. The control logic **204** can include circuitry, e.g., an integrated circuit integrating multiple logics, circuits, and/or components.

[0044] In some implementations, the control logic **204** includes a data register, an SRAM buffer, an address generator, a mode logic, and a state machine. The mode logic can be configured to determine whether there is a read or write operation and provide a result of the determination to the state machine.

[0045] During a write operation, the data register in the control logic **204** can register input data from the memory interface **202**, and the address generator in the control logic **204** can generate corresponding physical addresses to store the input data in specified memory cells of the memory cell array **210**. The address generator can be connected the X-decoder **206** and the Y-decoder that are controlled to select the specified memory cells through corresponding word lines and bit lines.

The SRAM buffer can retain the input data from the data register in its memory as long as power is being supplied. The state machine can process a write signal from the SRAM buffer and provide a control signal to a voltage generator that can provide a write voltage to the X-decoder **206** and/or the Y-decoder. The Y-decoder is configured to output the write voltage to the bit lines (BLs) for storing the input data in the specified memory cells.

[0046] During a read operation, the state machine can provide control signals to the voltage generator and the page buffer circuit **220**. The voltage generator can provide a read voltage to the X-decoder **206** and the Y-decoder for selecting a memory cell. A page buffer can sense a small power signal (e.g., a current signal) that represents a data bit (“1” or “0”) stored in the selected memory cell through a bit line **215** coupled to the page buffer and the selected memory cell. A sense amplifier can amplify the small power signal swing to recognizable logic levels so the data bit can be interpreted properly by logic inside or outside the memory device **200**. In some implementations, the page buffer circuit **220** and/or the data cache circuit **230** are included in the sense amplifier. The data buffer **240** (e.g., the data output buffer **244**) can receive the amplified voltage from the sensor amplifier and output the amplified power signal to the logic outside the memory device **200** through the memory interface **202**.

[0047] FIG. **3** illustrates an example of quantization error compensation for a multiply-and-accumulate (MAC) operation, according to some implementations. The upper half of FIG. **3** illustrates a MAC operation **300A** without compensation, whereas the lower half of FIG. **3** illustrates an MAC operation **300B** with compensation.

[0048] Both operations **300A** and **300B** involve an input floating point query vector **301** of (0.12, 0.26, 0.3) and a floating point weight vector (or floating point feature vector) **302** of (0.2, 0.3, 0.1). The task, which can be performed by IMC, is to compute a weighted sum of the elements of the floating point query vector **301** with the weights specified by the floating point weight vector **302**. When no error is incurred, the result of the task should be $0.12 \times 0.2 + 0.26 \times 0.3 + 0.3 \times 0.1 = 0.132$. However, when quantization is performed, there could be quantization error that lead to a different computing result.

[0049] For example, a 4-bit normalized quantization can be performed on the floating point query vector **301** and the floating point weight vector **302**. Under such quantization, the range $\{0, 1\}$ is equally divided into $2^{\text{sup.}4} = 16$ regions $\{0, 1/16\}$, $\{1/16, 2/16\}$, $\{2/16, 3/16\}$, \dots $\{15/16, 1\}$, which correspond to 16 integers 1, 2, 3, \dots 16, respectively. Because 0.12 falls within the region of $\{1/16, 2/16\}$, which corresponds to the integer of 2, 0.12 is quantized to 2. Similarly, 0.26 and 3 are quantized to 4 and 5, respectively. Accordingly, the floating point query vector **301** is quantized to an integer query vector **303** of (2, 4, 5). Likewise, the floating point weight vector **302** is quantized to an integer weight vector **304** of (3, 5, 2).

[0050] The integer representation of a floating point vector can lead to quantization error. For example, for an element that equals to 2 in the integer query vector **303**, any floating point value within the range of $\{1/16, 2/16\}$ are represented by the same integer. As such, when the integer representation is remapped to floating point values, the integer 2 may be rounded up to $2/16 = 0.125$, which is the upper boundary of the corresponding range, even though the exact floating point value is 0.12 before quantization. Following similar processes, the integer query vector **303** is remapped to a remapped floating point query vector **305** of (0.125, 0.25, and 0.312), whereas the integer weight vector **304** is remapped to a remapped floating point weight vector **306** of (0.187, 0.312, 0.125). In operation **300A** where no compensation is performed, the task using the remapped floating point query vector **305** and the remapped floating point weight vector **306** would return a result of $0.125 \times 0.187 + 0.25 \times 0.312 + 0.312 \times 0.125 = 0.140$. An error equal to $0.140 - 0.132 = 0.008$ has occurred due to quantization.

[0051] Different from operation **300A**, when a compensation vector ϵ is added to the remapped floating point query vector **305** prior to performing the task, then it is possible to reduce the quantization error. As an example, when $\epsilon = (0.003, 0, -0.072)$ is added to the remapped floating

point query vector **305**, a resulting compensated floating point query vector **305'** would equal (0.128, 0.25, 0.24). Performing the task using the compensated floating point query vector **305'** and the remapped floating point weight vector **306** would return a result of $0.128 \times 0.187 + 0.25 \times 0.312 + 0.24 \times 0.125 = 0.132$, with significant reduction of the error caused by quantization. Accordingly, by compensation with a carefully selected compensation vector $\&$, the accuracy of computing results can be significantly improved without also requiring that the quantization precision be increased. Furthermore, the compensation vector $\&$ can be a trainable or an untrainable compensation vector, which provides flexibility for IMC configuration of the AL/ML application.

[0052] FIG. 4 illustrates pseudo code of an example algorithm **400** for quantization error compensation, according to some implementations. The algorithm **400** can be performed by one or more processors, such as the host controller **122** and/or the processor **113** of the system **100**.

[0053] The algorithm **400** is described in the context of a computing task for determining cosine distances between an input query vector and a plurality of feature vectors so as to find the feature vector that has the highest cosine similarity to the input query vector. The computing task can be performed in part or in whole by the host controller **122**, the processor **113**, or the memory **116** of the system **100**. By finding the feature vector with the highest cosine similarity to the input query, the algorithm **400** can be practically adopted in AI or ML-based applications, such as query-based search (e.g., search for content with features closest to the features included in the query), brute force search (e.g., search for a specified target by exhausting all possibilities), and behavior-based content recommendation (e.g., provide a recommendation that matches a user's interest based on the user's or other users' past behavior). With the adoption of the algorithm **400** or the like in those applications, network techniques such as search engine optimization, route planning, metadata management, and network security can have improved efficiency, accuracy, and data integrity, and business and consumer experience in industries such as e-commerce, social media, and supply chain management can also be similarly improved.

[0054] As illustrated in FIG. 4, the algorithm **400** takes as inputs a query vector $X_{\text{sub.q}}$, a database vector D , and a quantization level N . The query vector $X_{\text{sub.q}}$ can be a floating point query vector representing one or more features of an input query, such as a search request, with the features extracted from user content, such as search criteria specified by the user in the search request. The database vector D can include one or more floating point feature vectors determined to be relevant to the features of the query, with each floating point feature vector representing, e.g., a candidate search result or group of search results within a database. The quantization level N can indicate a number of bits in the quantization process, which corresponds to the quantization precision.

[0055] Each of the one or more floating point feature vectors in the database vector D can correspond to an identifier (ID), such as an index, and the output of the algorithm **400** is the IDs of the one or more floating point feature vectors closest to the query vector $X_{\text{sub.q}}$, e.g., having the shortest cosine distances from the query vector $X_{\text{sub.q}}$. Put in a query-based search context, the return of the IDs of the closest floating point feature vectors can provide a user with search results that best match the criteria specified in the search request.

[0056] The algorithm **400** begins by quantization of the floating point query vector $X_{\text{sub.q}}$ at line 2 and quantization of the one or more floating point feature vectors $D_{\text{sub.i}}$ ($i=0, 1, 2 \dots$ being the index of the floating point feature vector), collectively represented as D , at line 3. Both quantizations have a precision according to the quantization level N . After quantization, the floating point query vector $X_{\text{sub.q}}$ becomes a quantized query vector $X_{\text{sub.q'}}$, and the one or more floating point feature vectors become one or more quantized feature vectors $D_{\text{sub.i'}}$, collectively represented as D' .

[0057] At line 4, the algorithm **400** determines a compensation vector $\&$ that minimizes the quantization error, represented as an error function. In some implementations, the error function is represented as a difference between i) a first result of performing the computing task after

compensation to the quantized query vector $X_{\text{sub}.q'}$, and ii) a second result that ideally would have been returned by the computing task should there be no floating point quantization. Here, the first result can be obtained by performing the computing task using $(X_{\text{sub}.q'} + E)$ and $D_{\text{sub}.i'}$ as inputs, whereas the second result can be obtained by performing the computing task using $X_{\text{sub}.q}$ and $D_{\text{sub}.i}$ as inputs.

[0058] In general, the cosine distance between two vectors can be computed by dividing the dot product of the two vectors by the product of the magnitudes of the two vectors. This is equivalent to multiplying the normalizations of the two vectors. Here, the normalization of a vector can be obtained by scaling the vector according to divisor that equals the magnitude of the vector.

Accordingly, for each floating point feature vector $D_{\text{sub}.i}$, the first result can be obtained as

$$[00001] \frac{\text{Math.}(X_q +) \cdot \text{Math.} D_{i'} \cdot \text{Math.}}{\text{Math.}(X_q +) \cdot \text{Math.} \cdot \text{Math.} \cdot \text{Math.} D_{i'} \cdot \text{Math.}},$$

and the second result can be obtained as

$$[00002] \frac{\text{Math.} X_q \cdot \text{Math.} D_i \cdot \text{Math.}}{\text{Math.} X_q \cdot \text{Math.} \cdot \text{Math.} \cdot \text{Math.} D_i \cdot \text{Math.}} \cdot$$

[0059] As such, for each floating point feature vector $D_{\text{sub}.i}$, an error function f can be obtained as

$$[00003] f = \frac{\text{Math.}(X_{q'} +) \cdot \text{Math.} D_{i'} \cdot \text{Math.}}{\text{Math.}(X_{q'} +) \cdot \text{Math.} \cdot \text{Math.} \cdot \text{Math.} D_{i'} \cdot \text{Math.}} - \frac{\text{Math.} X_q \cdot \text{Math.} D_i \cdot \text{Math.}}{\text{Math.} X_q \cdot \text{Math.} \cdot \text{Math.} \cdot \text{Math.} D_i \cdot \text{Math.}} \cdot$$

[0060] Reducing the magnitude of the error function f can correspond to reducing the computing error due to quantization. Accordingly, at line 4, a processor can be configured to determine a value ϵ of the compensation vector ϵ such that the magnitude of the error function f is minimized for each floating point feature vector $D_{\text{sub}.i}$. Here, the value $\epsilon_{\text{sub}.i}$ of the compensation vector ϵ collectively refers to the numeric values of the elements of the compensation vector ϵ . In other words, $\epsilon_{\text{sub}.i}$ can be determined based on

[00004]

$$\text{argmin } f() = \text{argmin } \left(\frac{\text{Math.}(X_{q'} +) \cdot \text{Math.} D_{i'} \cdot \text{Math.}}{\text{Math.}(X_{q'} +) \cdot \text{Math.} \cdot \text{Math.} \cdot \text{Math.} D_{i'} \cdot \text{Math.}} - \frac{\text{Math.} X_q \cdot \text{Math.} D_i \cdot \text{Math.}}{\text{Math.} X_q \cdot \text{Math.} \cdot \text{Math.} \cdot \text{Math.} D_i \cdot \text{Math.}} \right) \cdot$$

[0061] In some implementations, the determination of $\epsilon_{\text{sub}.i}$ is subject to a condition that the quantized query vector $X_{\text{sub}.q'}$ after compensation is approximately normalized with the floating point feature vector $D_{\text{sub}.i'}$:

$$[00005] \text{Math.} X_{q'}' + \text{Math.} \approx \text{Math.} D' \cdot \text{Math.} \cdot$$

[0062] In other words, the value of ϵ is determined such that a difference between $|X_{\text{sub}.q'} + \epsilon|$ and $|D'|$ is below a threshold value. This condition can be specific to the computing task of determining cosine distances. In other computing tasks, the determination of ϵ_i can be subject to different conditions.

[0063] With the one or more values $\epsilon_{\text{sub}.i}$ of the compensation vector ϵ determined, compensation can be made to the quantized query vector $X_{\text{sub}.q'}$ at line 5. The compensation can be the sum of the quantized query vector $X_{\text{sub}.q'}$ and the compensation vector ϵ having a value of $\epsilon_{\text{sub}.i}$.

[0064] At line 6 of the algorithm **400**, for each index i , the compensated quantized query vector $X_{\text{sub}.q'}$ and the corresponding floating point feature vector $D_{\text{sub}.i}$ are provide as inputs to perform the computing task of determining the cosine distance $\text{Dis}(i)$. As discussed above, the cosine distance between two vectors can be determined by normalizing the two vectors and obtaining the dot product of the two normalizations.

[0065] At line 7, the index i (or multiple indices) corresponding to the smallest cosine distance(s) $\text{Dis}(i)$ is identified. This can be achieved by, e.g., sorting all values of $\text{Dis}(i)$ in a descending order or an ascending order. At line 8, the identified index i (more multiple indices), or the sorted values of $\text{Dis}(i)$, is returned as an output of the computing task. In the example context of a query-based search application, the identified index can be used by the search application to locate the search result that best matches the search criteria specified in the user content.

[0066] Although the algorithm **400** described above is for determining cosine distances between a query vector and a plurality of feature vectors, some implementations similar to the algorithm **400**

can be adopted to determine other types of vector distances, such as Euclidean distances or Hamming distances. More generally, some implementations similar to the algorithm **400** can be adopted to perform other computing tasks according to the application and the user preference, e.g., the nature of the query or the underlying user content, the type of the input or output data, or available computing resources. The computing tasks can range from basic arithmetic functions, such as addition, subtraction, multiplication, and division, to more complex computing, such as matrix multiplication, fully connected layer computing, or k-nearest neighbors computing. For example, in the operation **300B** where the computing task is to find a weighted sum of the floating point query vector **301**, an algorithm can be adopted to determine an error function that corresponds to the computing task and to find the compensation vector $\varepsilon=(0.003, 0, -0.072)$ based on the error function. The operation **300B** can accordingly obtain the compensated floating point query vector **305'** and use it to perform the computing task.

[0067] FIG. **5** compares quantization error, e.g., difference between data distributions of a quantized query vector and a floating point query vector, with and without compensation, according to some implementations. The graph **500A** illustrates the quantization error without compensation, whereas the graph **500B** illustrates the quantization error with compensation according to, e.g., the algorithm **400** or a similar algorithm. In an ideal scenario, the non-quantized data samples and the quantized-dequantized data samples completely overlap, which suggests no quantization error has been introduced. In non-ideal scenarios, however, misalignment between two groups of data samples can be observed. The error function described above with reference to FIG. **4** can be determined based on the difference in data distribution between a compensated quantized query vector and the floating point query vector. By finding the compensation vector with a value that minimizes the magnitude of the error function, algorithms according to the implementations can reduce the quantization error.

[0068] Comparing the graphs **500A** and **500B**, it can be seen that the misalignment between bins representing non-quantized data samples and quantized-dequantized data samples is more visible in the graph **500A** than in the graph **500B**. This difference visibly shows the improvement to quantization and computing accuracy due to the compensation mechanisms according to some implementations.

[0069] FIGS. **6A** and **6B** together illustrate a flowchart of an example method **600** for performing a computing task based on a user content, according to some implementations. The method **600** can be performed, e.g., in part or in whole by the host controller **122**, the processor **113**, or the memory **116** of the system **100**. One or more operations of the method **600** can be similar to operations of algorithm **400**.

[0070] At **602**, one or more features are extracted from the user content. The one or more features can correspond to a search request indicated in the user content.

[0071] At **604**, the one or more features are converted to a floating point query vector, which can be similar to the query vector $X_{sub.q}$.

[0072] At **606**, the floating point query vector is quantized to obtain a quantized query vector, which can be similar to the quantized query vector $X_{sub.q'}$.

[0073] At **608**, a database vector is obtained based on the one or more features. The database vector, which can be similar to feature data set D , can include one or more floating point feature vectors.

[0074] At **610**, a compensation vector based on a data distribution of the floating point query vector is obtained.

[0075] At **612**, the one or more floating point feature vectors are quantized to obtain one or more quantized feature vectors, which can be similar to the quantized feature vectors D' .

[0076] At **614**, an error function is determined based on a difference between data distributions of i) the quantized query vector compensated with the compensation vector, and ii) the floating point query vector. The error function can be similar to the error function f .

[0077] At **616**, based on the error function, one or more values of the compensation vector corresponding to the one or more quantized feature vectors are determined. The determination of the one or more values of the compensation vector can be similar to the determination of E_i .

[0078] At **618**, the one or more quantized query vectors and the one or more values of the compensation vector are combined to obtain one or more compensated query vectors.

[0079] At **620**, the computing task is performed using the one or more compensated query vectors and the one or more quantized feature vectors to obtain an output.

[0080] At **622**, a response to the user content is generated based on the output. The response can indicate, e.g., a result to the search request.

[0081] The disclosed and other examples can be implemented as one or more computer program products, for example, one or more modules of computer program instructions encoded on a computer readable medium for execution by, or to control the operation of, data processing apparatus. The computer readable medium can be a machine-readable storage device, a machine-readable storage substrate, a memory device, or a combination of one or more of these. The term “data processing apparatus” encompasses all apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, or multiple processors or computers. The apparatus can include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, or a combination of one or more of them.

[0082] A system may encompass all apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, or multiple processors or computers. A system can include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, or a combination of one or more of them.

[0083] A computer program (also known as a program, software, software application, script, or code) can be written in any form of programming language, including compiled or interpreted languages, and it can be deployed in any form, including as a standalone program or as a module, component, subroutine, or other unit suitable for use in a computing environment. A computer program does not necessarily correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, sub programs, or portions of code). A computer program can be deployed for execution on one computer or on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communications network.

[0084] The processes and logic flows described in this document can be performed by one or more programmable processors executing one or more computer programs to perform the functions described herein. The processes and logic flows can also be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit).

[0085] Processors suitable for the execution of a computer program include, by way of example, both general and special purpose microprocessors, and any one or more processors of any kind of digital computer. Generally, a processor will receive instructions and data from a read only memory or a random access memory or both. The essential elements of a computer can include a processor for performing instructions and one or more memory devices for storing instructions and data. Generally, a computer can also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto optical disks, or optical disks. However, a computer need not have such devices. Computer readable media

suitable for storing computer program instructions and data can include all forms of nonvolatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

[0086] While this document may describe many specifics, these should not be construed as limitations on the scope of an invention that is claimed or of what may be claimed, but rather as descriptions of features specific to particular embodiments. Certain features that are described in this document in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable sub-combination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination in some cases can be excised from the combination, and the claimed combination may be directed to a sub-combination or a variation of a sub-combination. Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results.

[0087] Only a few examples and implementations are disclosed. Variations, modifications, and enhancements to the described examples and implementations and other implementations can be made based on what is disclosed.

Claims

1. A method for performing a computing task based on a user content, the method comprising: extracting one or more features from the user content, the one or more features corresponding to a search request indicated in the user content; converting the one or more features to a floating point query vector; quantizing the floating point query vector to obtain a quantized query vector; obtaining a database vector based on the one or more features, the database vector comprising one or more floating point feature vectors; determining a compensation vector based on a data distribution of the floating point query vector; quantizing the one or more floating point feature vectors to obtain one or more quantized feature vectors; determining an error function based on a difference between data distributions of i) the quantized query vector compensated with the compensation vector, and ii) the floating point query vector; determining, based on the error function, one or more values of the compensation vector corresponding to the one or more quantized feature vectors; combining the one or more quantized query vectors and the one or more values of the compensation vector to obtain one or more compensated query vectors; and performing the computing task using the one or more compensated query vectors and the one or more quantized feature vectors to obtain an output.
2. The method of claim 1, wherein determining the one or more values of the compensation vector comprises: determining the one or more values of the compensation vector such that a difference in magnitude between the quantized query vector compensated with the compensation vector and each of the one or more quantized feature vectors, is below a known threshold value.
3. The method of claim 1, wherein the computing task comprises performing at least one of a multiply-and-accumulate operation, general matrix multiplication (GeMM), fully connected layer computing, or k-nearest neighbors computing.
4. The method of claim 1, wherein the computing task comprises computing a plurality of vector distances.
5. The method of claim 4, wherein the plurality of vector distances comprises at least one of: a plurality of cosine similarity distances, a plurality of Euclidean distances, or a plurality of Hamming distances.

6. The method of claim 4, further comprising generating a response to the user content based on the output, wherein generating the response comprises sorting the plurality of vector distances in a descending order or an ascending order.
7. The method of claim 1, wherein the user content comprises at least one of: graphical information, textual information, geographical information, or temporal information.
8. The method of claim 1, further comprising receiving the user content from at least one of: a text-based search engine, a graph-based search engine, a brute force search engine, or a behavior-based content recommendation system.
9. The method of claim 1, wherein performing the computing task comprises performing the computing task using an in-memory computing (IMC) circuit, and wherein the IMC circuit comprises a plurality of memory cells comprising at least one of NAND flash cells, NOR flash cells, phase change memory (PCM), Magnetoresistive random-access memory (MRAM), Ferroelectric random-access memory (FeRAM), or Spin-Transfer-Torque random-access memory (STT-RAM).
10. A computing system comprising: one or more processors; and a computing circuit coupled to the one or more processors, wherein the one or more processors are configured to execute instructions to perform operations comprising: extracting one or more features from user content, the one or more features corresponding to a search request indicated in the user content; converting the one or more features to a floating point query vector; quantizing the floating point query vector to obtain a quantized query vector; obtaining a database vector based on the one or more features, the database vector comprising one or more floating point feature vectors; determining a compensation vector based on a data distribution of the floating point query vector; quantizing the one or more floating point feature vectors to obtain one or more quantized feature vectors; determining an error function based on a difference between data distributions of i) the quantized query vector compensated with the compensation vector, and ii) the floating point query vector; determining, based on the error function, one or more values of the compensation vector corresponding to the one or more quantized feature vectors; combining the one or more quantized query vectors and the one or more values of the compensation vector to obtain one or more compensated query vectors; and performing, by the computing circuit, the computing task using the one or more compensated query vectors and the one or more quantized feature vectors to obtain an output.
11. The computing system of claim 10, wherein determining the one or more values of the compensation vector comprises: determining the one or more values of the compensation vector such that a difference in magnitude between the quantized query vector compensated with the compensation vector and each of the one or more quantized feature vectors, is below a known threshold value.
12. The computing system of claim 10, wherein the computing task comprises performing at least one of: a multiply-and-accumulate operation, general matrix multiplication (GeMM), fully connected layer computing, or k-nearest neighbors computing.
13. The method of claim 12, wherein the computing task comprises computing a plurality of vector distances, and wherein the plurality of vector distances comprises at least one of: a plurality of cosine similarity distances, a plurality of Euclidean distances, or a plurality of Hamming distances.
14. The method of claim 12, the operations further comprising generating a response to the user content based on the output, wherein generating the response comprises sorting the plurality of vector distances in a descending order or an ascending order.
15. The computing system of claim 10, wherein the user content comprises at least one of: graphical information, textual information, geographical information, or temporal information.
16. The computing system of claim 10, the operations further comprising receiving the user content from at least one of: a text-based search engine, a graph-based search engine, a brute force search engine, or a recommendation system.

17. The computing system of claim 10, wherein the computing circuit comprises an in-memory computing (IMC) circuit, wherein the IMC circuit comprises a plurality of memory cells comprising at least one of: NAND flash cells, NOR flash cells, phase change memory (PCM), Magnetoresistive random-access memory (MRAM), Ferroelectric random-access memory (FeRAM), or Spin-Transfer-Torque random-access memory (STT-RAM).

18. A non-transitory computer-readable medium storing program instructions that, when executed, cause one or more processors to perform operations comprising: extracting one or more features from user content, the one or more features corresponding to a search request indicated in the user content; converting the one or more features to a floating point query vector; quantizing the floating point query vector to obtain a quantized query vector; obtaining a database vector based on the one or more features, the database vector comprising one or more floating point feature vectors; determining a compensation vector based on a data distribution of the floating point query vector; quantizing the one or more floating point feature vectors to obtain one or more quantized feature vectors; determining an error function based on a difference between data distributions of i) the quantized query vector compensated with the compensation vector, and ii) the floating point query vector; determining, based on the error function, one or more values of the compensation vector corresponding to the one or more quantized feature vectors; combining the one or more quantized query vectors and the one or more values of the compensation vector to obtain one or more compensated query vectors; and performing the computing task using the one or more compensated query vectors and the one or more quantized feature vectors to obtain an output.

19. The non-transitory computer-readable medium of claim 18, wherein determining the one or more values of the compensation vector comprises: determining the one or more values of the compensation vector such that a difference in magnitude between the quantized query vector compensated with the compensation vector and each of the one or more quantized feature vectors, is below a known threshold value.

20. The non-transitory computer-readable medium of claim 18, wherein the computing task comprises computing a plurality of vector distances.
