

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265273

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

RAHUT; Anirban et al.

SEARCH RESULT REPLICATION MANAGEMENT IN A SEARCH HEAD CLUSTER

Abstract

Systems and methods for search result replication in a search head cluster of a data aggregation and analysis system. An example method may include receiving, by a search head leader of a search head cluster including multiple search heads, from a first search head of the plurality of search heads, a search result in response to a search query. The search head leader parses a registry comprising a set of replicas of the search result in the search head cluster to determine a replication count corresponding to a number of replicas of the search result. A determination is made that the replication count is greater than a target replication count. Based on the determination, a least-recently-accessed replica from the set of replicas is identified and a deletion of the least-recently-accessed replica is initiated.

Inventors: RAHUT; Anirban (Santa Clara, CA), Vasan; Sundar (San Francisco, CA)

Applicant: Splunk Inc. (San Francisco, CA)

Family ID: 1000008575057

Appl. No.: 19/184991

Filed: April 21, 2025

Related U.S. Application Data

parent US continuation 18201042 20230523 parent-grant-document US 12282497 child US 19184991

parent US continuation 16159893 20181015 parent-grant-document US 11704341 child US 18201042

parent US continuation 14449069 20140731 parent-grant-document US 10133806 child US 16159893

Publication Classification

Int. Cl.: **G06F16/28** (20190101); **G06F16/27** (20190101); **G06F16/951** (20190101);
G06F16/9538 (20190101)

U.S. Cl.:

CPC **G06F16/285** (20190101); **G06F16/27** (20190101); **G06F16/951** (20190101);
G06F16/9538 (20190101);

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS [0001] This application is a continuation of U.S. patent application Ser. No. 18/201,042, filed on May 23, 2023, which itself is a continuation of U.S. Pat. No. 11,704,341, filed on Oct. 15, 2018, which is a continuation of U.S. Pat. No. 10,133,806, filed on Jul. 31, 2014. The entire contents are hereby incorporated by reference herein.

TECHNICAL FIELD

[0002] The present disclosure is generally related to data aggregation and analysis systems, and is more specifically related to search result replication in a search head cluster.

BACKGROUND

[0003] Modern data centers often comprise thousands of hosts that operate collectively to service requests from even larger numbers of remote clients. During operation, components of these data centers can produce significant volumes of machine-generated data. The unstructured nature of much of this data has made it challenging to perform indexing and searching operations because of the difficulty of applying semantic meaning to unstructured data. As the number of hosts and clients associated with a data center continues to grow, processing large volumes of machine-generated data in an intelligent manner and effectively presenting the results of such processing continues to be a priority.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] The present disclosure is illustrated by way of examples, and not by way of limitation, and may be more fully understood with references to the following detailed description when considered in connection with the figures, in which:

[0005] FIG. 1 presents a block diagram of an event-processing system for performing search result replication in a search head cluster, in accordance with the disclosed embodiments.

[0006] FIG. 2 depicts a flow diagram of an example method for search head leader election, in accordance with one or more aspects of the present disclosure.

[0007] FIG. 3A presents a block diagram illustrating an exemplary search head leader to perform search result replication in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0008] FIG. 3B presents a block diagram illustrating an exemplary search head follower to receive search result replication in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0009] FIG. 4 depicts a flow diagram of an example method for search result replication in a search

head cluster, in accordance with one or more aspects of the present disclosure.

[0010] FIG. 5 schematically illustrates an example acknowledgement communication flow for performing search result replication in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0011] FIG. 6 depicts a flow diagram of an example method for a search head failure process in association with search replication in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0012] FIG. 7 depicts a flow diagram of an example method for removing excess copies of search results in association with search replication in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0013] FIG. 8 depicts a flow diagram of an example method for search result proxying utilizing replicated search results in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0014] FIG. 9 presents a block diagram of an event-processing system in accordance with the disclosed embodiments.

[0015] FIG. 10 presents a flowchart illustrating how indexers process, index, and store data received from forwarders in accordance with the disclosed embodiments.

[0016] FIG. 11 presents a flowchart illustrating how a search head and indexers perform a search query in accordance with the disclosed embodiments.

[0017] FIG. 12 presents a block diagram of a system for processing search requests that uses extraction rules for field values in accordance with the disclosed embodiments.

[0018] FIG. 13 illustrates an exemplary search query received from a client and executed by search peers in accordance with the disclosed embodiments.

[0019] FIG. 14A illustrates a search screen in accordance with the disclosed embodiments.

[0020] FIG. 14B illustrates a data summary dialog that enables a user to select various data sources in accordance with the disclosed embodiments.

[0021] FIG. 15A illustrates a key indicators view in accordance with the disclosed embodiments.

[0022] FIG. 15B illustrates an incident review dashboard in accordance with the disclosed embodiments.

[0023] FIG. 15C illustrates a proactive monitoring tree in accordance with the disclosed embodiments.

[0024] FIG. 15D illustrates a screen displaying both log data and performance data in accordance with the disclosed embodiments.

[0025] FIG. 16 depicts a block diagram of an example computing device operating in accordance with one or more aspects of the present disclosure.

DETAILED DESCRIPTION

[0026] Disclosed herein are systems and methods for search result replication in a search head cluster. The search head cluster may include one or more search heads of a data aggregation and analysis system grouped together to form a cluster. A search head may refer to a component of the data aggregation and analysis system that is responsible for a reduce phase of a map-reduce search performed by the data aggregation and analysis system, as further described below. A search head may communicate with one or more indexers that are responsible for a map phase of the map-reduce search, as further described below.

[0027] An example data aggregation and analysis system may aggregate heterogeneous machine-generated data received from various sources, including servers, databases, applications, networks, etc. The aggregated source data may comprise a plurality of events. An event may be represented by a data structure that is associated with a certain point in time and comprises a portion of raw machine data (i.e., machine-generated data). The system may be configured to perform real-time indexing of the source data and to execute real-time, scheduled, or historic searches on the source data. A search query may comprise one or more search terms specifying the search criteria. Search

terms may include keywords, phrases, Boolean expressions, regular expressions, field names, name-value pairs, etc. The search criteria may comprise a filter specifying relative or absolute time values, to limit the scope of the search by a specific time value or a specific time range. The search query may also process retrieved data, such as by determining aggregates or calculating statistics or deriving values for performance metrics.

[0028] The search result may be stored on a local disk of the search head that executed the search query. The search result may also be referred to as an artifact, a search directory (e.g., referring to the location where the search result is stored), a dispatch directory (e.g., referring to the location where the search result is stored), a resulting dataset, or a search identifier (ID) (e.g., referring to the identifier of the search result), to name a few examples.

[0029] The example data aggregation and analysis system may be employed to replicate search results across search heads of a search head cluster implemented by the data aggregation and analysis system associated with a distributed computer system (e.g., an enterprise system comprising a plurality of computer systems and peripheral devices interconnected by a plurality of networks). Accordingly, implementations of the disclosure provide an effective mechanism for providing high availability, data recovery, and network communication efficiency for a data aggregation and analysis system by replicating search results generated by the system.

[0030] Various aspects of the methods and systems are described herein by way of examples, rather than by way of limitation. The methods described herein may be implemented by hardware (e.g., general purpose and/or specialized processing devices, and/or other devices and associated circuitry), software (e.g., instructions executable by a processing device), or a combination thereof.

[0031] FIG. 1 presents a block diagram of an exemplary event-processing system **100** (e.g., a data aggregation and analysis system) (hereinafter referred to as “system **100**”) for replicating search results in a search head cluster, in accordance with one or more aspects of the present disclosure. In one implementation, system **100** is similar to the SPLUNK® ENTERPRISE system produced by Splunk Inc. of San Francisco, California, which is described in more details herein below). System **100** is also described in further detail below with respect to FIG. 9.

[0032] System **100** includes one or more indexers **102** that store, process, and/or perform operations on data obtained from a variety of different sources, wherein each indexer operates on data contained in a specific data store **103**. These indexers **102** can comprise separate computer systems in a data center, or may alternatively comprise separate processes executing on various computer systems in a data center.

[0033] During operation, the indexers **102** receive the collected data, which is distributed across different indexers to facilitate parallel processing. This parallel processing can take place at data ingestion time, because multiple indexers **102** can process the incoming data in parallel. The parallel processing can also take place at search time, because multiple indexers **102** can search through their respective portions of the collected data in parallel.

[0034] Each indexer **102** is responsible for storing and searching a different subset datastore **103** of all of the events contained collectively within the system. As discussed in more detail below with respect to FIG. 9, an “event” may comprise a portion of performance data and/or diagnostic information that is generated by a computer system and is correlated with a specific point in time. By distributing events among the indexers **102** and data stores **103**, the indexers **102** can analyze events for a query in parallel, for example using map-reduce techniques (described further below with respect to FIG. 9), wherein each indexer **102** returns partial responses for a subset of events to a search head **104** that combines the results to produce an answer for the query. By storing events in buckets for specific time ranges, an indexer **102** may further optimize searching by looking only in buckets for time ranges that are relevant to a query.

[0035] The indexers **102** may be communicably coupled to one or more search heads **104A-B** of a search head cluster **120**. A search head **104A-B** may comprise a server device accessible by a number of client devices (not shown) over a network. Search head **104A-B** is a component that is

responsible for receiving and processing a search query. The search query may be sent from a client (e.g., a browser communicating with system **100**). The search head **104A-B** may analyze the received search query and determine which portions of the query may be delegated to indexers **102** and what portions should be executed locally by the search head **104A-B**. The search head **104A-B** may be responsible for a reduce phase of a map-reduce search by combining partial results and/or events received from indexers **102** to produce a final search result for the query. This final search result can comprise different types of data depending on what the query is asking for. Further details on the responsibilities and operations of search head **104A-B** are discussed below with respect to FIG. **11**.

[0036] Although the search head **104A-B** is described as having the duties of both performing searches to produce results and making previously produced search results available to clients, these duties could be separated, with the entities storing the results of the search and making them available to clients doing so according to the methods of the present disclosure.

[0037] In some implementations, the system **100** may horizontally scale the search capability (e.g., increase the number of searches than can be run) of system **100** by creating a pool of search heads **104A-B** referred to as a search head cluster **120**. Within the search head cluster **120**, each search head **104A-B** is communicably coupled to each indexer **102**. The search heads **104A-B** may be located at one or more sites, each site being remote from the other site, e.g., in different geographic locations. In one embodiment, all of the search heads **104A-B** of the search head cluster **120** are located at a single site. In another embodiment, at least two of the search heads **104A-B** are located at a single site and at least one other of the search heads **104A-B** is located at a different site. In another embodiment, each search head **104A-B** is located at a different site.

[0038] It may be advantageous for the search heads **104A-B** to interact with each other as a single server cluster, referred to as the search head cluster **120**. At a search head cluster **120**, instructions to perform a task received by one of the search heads **104A-B** may be performed by any of the search heads **104A-B** according to the load on each of the search heads **104A-B**.

[0039] The search head cluster **120** may be operated in a symmetric mode in which each search head **104A-B** has an equal role or in an asymmetric mode in which one of the search heads **104A** is designated the “leader” or “captain” (referred to herein as “search head leader”) and has authority to issue commands to other search heads **104B** (referred to herein as “search head followers”) and make decisions regarding the system. For example, each search head **104A-B** may forward instructions from client devices to perform tasks according to a schedule to the search head leader **104A**. In turn, the search head leader **104A** may schedule the tasks to be performed by the various search head followers **104B** by issuing commands to the various search head followers **104B** to perform the tasks at specified times. Because, as described below, any of the search heads **104A-B** may be designated the leader, each of the search heads **104A-B** includes a scheduling module (not shown) to perform this task in the event the search head **104A-B** is elected the leader.

[0040] In the asymmetric mode, each search head **104A-B** is either in a leader state, a follower state, or a candidate state. In the leader state, the search head **104A** acts as the leader and makes decisions regarding the cluster **120** and issues commands to other search head followers **104B**. In the follower state, the search head follower **104B** receives commands from the search head leader **104A** and follows the commands. A search head **104A-B** may enter the candidate state during a leader election process as described further below. Each of the search heads **104A-B** includes an election module (not shown) to perform a leader election process. The leader election process may be, for example, a consensus election process that uses the Raft consensus protocol as discussed in more detail below. The Raft consensus protocol is further described in “In Search of Understandable Consensus Algorithm (Extended Version)” by Diego Ongaro and John Ousterhout, Stanford University, 2014, which is hereby incorporated herein by reference in its entirety for all purposes.

[0041] Other consensus algorithms or leader selection protocols may also be utilized in

implementations of the disclosure. The purpose of the Raft protocol or other consensus election protocols is to ensure that there is always a leader search head for a search head cluster. In the case that a search head leader fails or is disconnected from the other search heads, the Raft protocol or other consensus election protocols may be utilized to elect a new search head leader.

[0042] With respect to the leader election process, during operation of the distributed system, time may be partitioned into terms. During each term, the search heads **104A-B** attempt to elect a leader and, if a leader is elected, operate in the asymmetric mode. A term may begin during system initiation or when a search head **104A** acting as the leader fails or is otherwise non-responsive (such as a network communication failure). Each search head **104A-B** may store a numeric value identifying the current term (e.g., 1, 2, . . . N, etc.). This current term value may be stored in non-volatile memory such that when a search head **104A-B** fails and is restarted, the search head **104A-B** may access the current term value. During each term, each search head **104A-B** may cast a single vote to elect one of the search heads **104A-B** as a leader. Thus, each search head **104A-B** may store, along with the current term value, an identifier of a search head **104A-B** for which the search head **104A-B** has voted to be a leader during the current term. The identifier may be stored in association with the current term value in non-volatile memory.

[0043] FIG. 2 presents a flowchart illustrating a method **200** of how a search head, such as search head **104A-B** described with respect to FIG. 1, operates in a distributed system to elect a leader. In one implementation, method **200** utilizes the Raft consensus protocol for the leader selection. When a search head is initialized, the search head starts in the follower state in block **210**. The search head may be initialized when the distributed system is initialized. The search head may also be initialized because it has joined an existing distributed system or it failed and has been restarted. If a leader has been elected, the leader sends out periodic heartbeat messages to maintain its authority. The heartbeat messages may be transmitted approximately every 150 milliseconds. The heartbeat messages serve to inform the other search heads that the leader is operational. Accordingly, heartbeat messages may, in some circumstances, not be sent to a search head if the leader is sending other messages to the search head as these other messages serve to inform the search head that the leader is operational.

[0044] Thus, in the follower state, the search head expects to receive heartbeat messages from a leader on a periodic basis. If a period of time passes without the search head receiving a heartbeat message from a leader, referred to as a heartbeat timeout, the search head may begin an election process to elect a new leader. Thus, at block **220**, the search head determines whether a heartbeat message has been received. If a heartbeat message has been received, the search head returns to block **220** to await another heartbeat message. If a heartbeat message has not been received, the search head continues to block **230** where it determines whether a heartbeat timeout has occurred. The heartbeat timeout may be randomly selected by the search head between a first value and a second value. For example, the heartbeat timeout may be selected by the search head as any value between 100 and 500 milliseconds. By randomly selecting a heartbeat timeout, each search head decides to begin the election process at a different time, preventing all of the search heads simultaneously transitioning to candidate state and requesting votes (as described below).

[0045] If the search head determines, at block **230**, that a heartbeat timeout has not occurred, the search head returns to block **220** and awaits a heartbeat message. If the search head determines that a heartbeat timeout has occurred, the search head continues to block **240** where it enters the candidate state.

[0046] Once in the candidate state, the search head increments the current term value (e.g., by 1) stored in memory (at block **250**), votes for itself (at block **260**), and requests votes from the other search heads in the distributed system (at block **270**). The search head may vote for itself by associating its own identifier with the updated current term value in non-volatile memory. The search head may request votes from the other search heads by transmitting messages to the other search heads including, among other things, the incremented current term value and the identifier

of the requesting search head.

[0047] Each search head, upon receiving a request for a vote, may respond to the requesting search head. If the receiving search head has already not voted for another search head in the term indicated by the request, the receiving search head may vote for the requesting search head by storing the requesting search head's identifier in memory in association with the term value and sending a response to the requesting search head indicating that the receiving search head has cast its vote for the requesting search head. If the receiving search head has already voted for another search head, as indicated by another search head's identifier stored in the receiving search head's memory in association with the term value that matches the term value in the request, the search head may send a response to the requesting search head indicating that the request is denied.

[0048] At block **280**, the search head determines the election results. If the search head receives a majority of votes, the election is won and the search head continues to block **290** where it enters into the leader state. The search head may determine that it has received a majority of the votes when the number of responses from other search heads (plus one for its own vote for itself) is greater than half the number of search heads in the distributed system. Once in the leader state, the search head sends out heartbeat messages to the other search heads, thereby indicating that it has taken on the leadership role. The other search heads may respond by sending heartbeat acknowledgement signals back to the leader. In one embodiment, if the leader does not receive heartbeat acknowledgment signals from at least a majority of the search heads in the distributed system, the search head abdicates its leadership role by transitioning to the follower state, returning to block **210**.

[0049] If the search head receives a heartbeat message from another search head indicating that the other search head has transitioned to the leader state, the election is lost and the search head returns to block **210** where it enters into the follower state.

[0050] If a period of time elapses without the search head receiving a majority of votes or receiving a heartbeat message indicating that another search head has become the leader, referred to as an election timeout, the search head returns to block **250**, where another term is begun and the election process begins again. An election timeout may occur when multiple search heads have transitioned to the candidate state and votes are split such that none of the search heads receive a majority of the votes. The election timeout may be the same as the heartbeat timeout or a different value. In one embodiment, the election timeout is randomly selected between 100 and 500 milliseconds.

[0051] Referring back to FIG. 1, when a search head cluster **120** is implemented by system **100**, search results stored locally on one search head **104A-B** may be not available to a client connected to another search head **104A-B** of the search head cluster **120**. Furthermore, if a search head fails, any search results stored locally at the search head **104A-B** may be lost. In such a case, the search may have to be re-executed, resulting in runtime inefficiencies. Thus, in implementations of the disclosure, the search head leader **104A** is responsible for search result replication among the search heads **104A-B** of the search head cluster **120** in order to provide high availability of the search results stored at the search heads **104A-B** of the search head cluster **120**. In other implementations, the search head leader **104A** is also responsible for proxying of search result requests to search heads to store a copy of the requested search results. The search result replication and proxying of implementations of the disclosure are described below in further detail with respect to FIGS. 3A-8.

[0052] FIG. 3A presents a block diagram of an exemplary search head leader **104A**, in accordance with one or more aspects of the present disclosure. Search head leader **104A** may be the same as search head leader **104A** described above with respect to FIG. 1, and selected per method **200** described above with respect to FIG. 2. As discussed above, if search head leader **104A** fails, any of search head followers **104B** may be elected as new search head leader **104A**, and thus may also include the components detailed for search head leader **104A** in FIG. 3A. In one implementation, search head leader **104A** may comprise a computing device, such as a server device, and/or

software or firmware executing on a computing device.

[0053] Referring to FIG. 3A, search head leader **104A** may include a replication component **310**. Replication component **310** is responsible for managing and performing search result replication among the search heads of the search head cluster. Replication component **310** may maintain and reference a cluster search results registry **330** and a replication tracking registry **340** as part of a replication process performed by replication component **310**. The cluster search results registry **330** may be any type of data store, such as table or database, which maps all of the search results stored in the search heads of the search head cluster. The cluster search results registry **330** may include, but is not limited to, for each search result, a search result ID (SID), a source search node that produced the search result, one or more search nodes where the search result is replicated, a number of replications, and metadata associated with the search result (e.g., how long the search ran, what was the search string, when did the search start, when did the search end, interesting fields (e.g., fields specified in search terms of the search), execution statistics of search, etc.). When a search is initially completed by a search node, the search node provides the search result information detailed above to the search head leader **104A** for inclusion in the cluster search results registry **330**.

[0054] In one implementation, the cluster search results registry **330** is maintained in non-persistent memory of the search head leader **104A**. If the search head leader **104A** fails and a new search head is elected as leader, the new search head can request and/or receive search result information from the other search head members in the cluster in order to re-generate the cluster search results registry **330**. In other implementations, the cluster search results registry **330** is maintained in persistent storage of the search head leader **104**.

[0055] In some implementations, the search head leader **104A** provides data of the cluster search results registry **330** to the other search heads (i.e., search head followers) in the search head cluster in order for the other search heads to maintain their own local copy of the cluster search results registry **330**. FIG. 3B presents a block diagram of an exemplary search head follower **104B**, in accordance with one or more aspects of the present disclosure. Search head follower **104B** may be the same as search head followers **104B** described above with respect to FIG. 1. In one implementation, search head follower **104B** may comprise a computing device, such as a server device, and/or software or firmware executing on a computing device.

[0056] As shown in FIG. 3B, the search head follower **104B** may include a local cluster search results registry **360** generated based on data provided from the search head leader **104A** describing the state of the primary cluster search results registry **330** maintained by the search head leader **104A**. In one implementation, the search head leader **104A** periodically sends data representing a delta of the cluster search results registry **330** since a previous status update was sent by the search head leader **104A**. Providing a delta reduces the amount of data to be sent over the network by the search head leader **104A** to the search head followers **104A**. In addition, in some implementations, a search head follower **104B** may utilize its local cluster search results registry **360** to re-generate the primary cluster search results registry **330** if the search head follower **104B** is elected as a new search head leader **104A**. In addition, in some implementations, the search head follower **104B** includes a proxying component **320**, similar to the search head leader **104A**. Details of the proxying component **320** are described further below with respect to FIG. 8.

[0057] Referring back to FIG. 3A, the replication tracking registry **340** may be any type of data store, such as table or database, which tracks in-progress replications occurring between search heads in the search head cluster. The replication tracking registry **340** may include, but is not limited to, an SID, a source search head originating the search result, a target search head receiving the search result replication, and a time when the replication started. The replication tracking registry **340** may be maintained in non-persistent memory of the search head leader **104A**. In other implementations, the cluster search results registry **330** is maintained in persistent storage of the search head leader **104**.

[0058] Each search head in a cluster, including the search head leader **104A** and all search head followers **104B**, may maintain a data store of search results **350**. The search results **350** may include both local search results **352** performed by the respective search head and replicated search results **354** performed by other search heads in the cluster and replicated to the respective search head. Although illustrated as separate components, the local search results **352** and the replicated search results **354** may be stored together in a single data store.

[0059] Each search head **104A-B** may also store a numeric value identifying a default (or target) replication count that configures the number of replicas of the search results to be maintained in a cluster. In one implementation, the default replication count is configured by an administrator of the data aggregation and analysis system and stored in a configuration file propagated to and maintained by each search head in a cluster. This default replication count may be stored in non-volatile memory such that when a search head **104A-B** fails and is restarted, the search head **104A-B** may access the default replication count.

[0060] The replication component **310** of search head leader **104A** may manage a replication process for the search head cluster. FIGS. **4-7** provide further detail of the replication process performed by replication component **310**. FIGS. **4-8** depict flow diagrams of example methods **400-800** for search result replication and proxying in a search head cluster. Methods **400-800** and/or each of their respective individual functions, routines, subroutines, or operations may be performed by one or more general purpose and/or specialized processing devices. Two or more functions, routines, subroutines, or operations of methods **400-800** may be performed in parallel or in an order that may differ from the order described above. In certain implementations, one or more of methods **400-800** may be performed by a single processing thread. Alternatively, methods **400-800** may be performed by two or more processing threads, each thread executing one or more individual functions, routines, subroutines, or operations of the respective method. In an illustrative example, the processing threads implementing methods **400-800** may be synchronized (e.g., using semaphores, critical sections, and/or other thread synchronization mechanisms). Alternatively, the processing threads implementing methods **400-800** may be executed asynchronously with respect to each other. In an illustrative example, methods **400-800** may be performed by an example computing device **1000** described herein below with reference to FIG. **16**. In another illustrative example, methods **400-800** may be performed by a distributed computer system comprising two or more example computing devices **1000**.

[0061] FIG. **4** depicts a flow diagram of an example method **400** for search result replication in a search head cluster, in accordance with one or more aspects of the present disclosure. Referring to FIG. **4**, at block **425**, the computer system implementing the method may receive a search result from a completed search from one of the search heads in a search head cluster. The search result may be the product of the execution of a search query by one or more indexers and aggregated by the sending search head. In an illustrative example, the search query may represent a real-time search (e.g., may repeatedly be executed by a certain process or thread in an indefinite loop which may be interrupted by occurrences of certain terminating conditions). In another illustrative example, the search query may represent a scheduled search (e.g., may be executed according to a certain schedule). The search result may correspond to a search scheduled by the search head leader and forwarded to a search head follower for execution, or a search managed by the search head leader. The search result provided to the search head leader may include an SID of the search result and metadata associated with the search result, such as how long the search ran, the search string, when the search started, when the search ended, interesting fields of the search, execution statistics of the search, and so on.

[0062] Responsive to receiving the search result, at block **430**, the search head leader creates a new entry for the search result in a cluster search results registry maintained by the search head leader. The cluster search results registry may be the same as the cluster search results registry **330** described above with respect to FIG. **3A**. At block **435**, the new search result entry is monitored for

compliance with a configured replication count of the search head cluster. The replication count dictates how many locations in a search head cluster a search result should be replicated.

[0063] In some implementations, search results that are in an intermediate or in-transit state are not replicated until a final set of results is received at the search head leader. In addition, the search results that are “ad-hoc” in nature may not be replicated by the search head leader. An “ad-hoc” search result may refer to a search result that is performed one time and stored by the search head for a limited period of time. Furthermore, the search head leader may replicate the search results regardless of any sharing permissions associated with the search results. For example, if a search result is configured as a private search result that is not shared with any users other than the originating user, then this search result may still be replicated in order to provide high availability of the search result within the search head cluster. In other implementations, all search results, regardless of intermediate, in-transit, ad-hoc, scheduled, and/or permissions state are replicated. In further implementations, formatting metadata defining how the search result should be presented (or defining a visualization for the search result) as specified by an originating search query user may be replicated with the search result.

[0064] If, at decision block **440**, the search result is in compliance (e.g., number of occurrences of search result in the cluster meets or exceeds the default replication count), then method **400** returns to block **435** to continue monitoring the search result entry. On the other hand, if it is determined, at decision block **440**, that the search result is not in compliance (e.g., number of occurrences of search result in the cluster is less than the default replication count), then method **400** proceeds to block **445**.

[0065] At block **445**, a target search head is selected from the search head cluster to receive a replication of the search result. In one implementation, the search head leader performs a candidate selection process in order to determine the target search head to receive the replication. The candidate selection process may comprise applying a heuristic including, but not limited to, a random selection of a target search head, a site-aware selection of a target search head, and/or a balanced selection of a target search head. The site-aware selection may refer to selecting a target search head that is located at a different site (e.g., geographically) than the source search head that originated the search result and/or the other search heads that store copies (i.e., replicas) of the search result. In some implementations, the configuration settings associated with the particular search results may dictate that such a site-awareness policy be applied to the search result, and may indicate specific geographic sites of preference for such replications. In another implementation, the search head leader may apply learning algorithms to determine on which site a search result is most popular (e.g., keep count of how many accesses have occurred for the search result on a particular site), and based on that learning, may perform bias replications of the search result to the determine popular site(s).

[0066] The balanced selection may refer to distributing replications across the search head cluster in a balanced manner. For example, a replication may be biased to a search head that has the smallest number of search results stored at the search head, or has the smallest amount of search result memory consumed.

[0067] Responsive to the selection of a target search head, at block **450**, a replication process is initiated from the source search head originating the search result to the target search head. FIG. 5 schematically illustrates an exemplary replication process communication flow between the search head leader **104A**, the source search head **510**, and the target search head **520**. The search head leader **104A** initiates a replication process by sending a request **530** to the source search head **510** to start replicating the search result to the target search head **530**. The source search head acknowledges **531** to the search head leader **104A** that the source search head **510** is starting the replication process and asks the search head leader **104A** to begin monitoring the replication process. In one implementation, the search head leader **104A** may create a new entry in the replication tracking registry **340** of FIG. 3A to utilize in monitoring the in-progress replication

process.

[0068] The source search head **510** also opens a channel with the target search head **510** via open channel negotiations **532**, **533**. The target search head **520**, upon receiving the request to open the channel for the replication process from the source search head **510**, acknowledges **534** to the search head leader **104A** that the target search head is beginning to participate in the replication process. The source search head **510** then begins sending files of the search result **535** to replicate in a compressed format to the target search head **520** over the established channel. In one implementation, the underlying mechanism for sending and compressing the search result may utilize any known protocol that sends data in an efficient format. After every block of data is sent **535**, the target search head **520** acknowledges **536** receipt of the data back to the source search head **510**.

[0069] At the end of sending files for the search result, the source search head **510** sends a signal indicating that the data transfer is done **537** to the target search head **520**. The source search head also acknowledges **538** back to the search head leader that the source search head **510** is done with the replication process. When the target search head **520** receives the done signal and all of the replicated data is pushed to search result storage (e.g., search results **350** of FIGS. **3A-3B**) of the target search head **520**, the target search head **520** acknowledges the done signal to the source search head **510** and acknowledges **540** to the search head leader **104A** that the target search head **520** has completed its portion of the replication process.

[0070] Utilizing the acknowledgements received from the source search head **510** and the target search head **520**, as tracked in the replication tracking registry **340**, the search head leader **104A** can coordinate and monitor the replication process for a search result. The search head leader **104A** can determine, based on this information, whether a replication has completed or failed. For example, the search head leader **104A** may utilize the tracked acknowledgements from the source search head **510** and the target search head **520**, as well as failure notifications from the source search head **510** and the target search head **520**, to determine when a replication has failed between two search heads. The search head leader **104A** can then remove (flush) the tracked replication from the replication tracking registry **340**, select another target search head to receive the search result replication, and re-start the replication process for the search result.

[0071] Referring back to FIG. **4**, after the replication process is determined to be complete, at block **455**, a replication count corresponding to the search result is incremented. In one implementation, the replication count corresponding to the search result is maintained in the cluster search result registry at the search head leader.

[0072] FIG. **6** depicts a flow diagram of an example method **600** for a search head failure process in associated with search replication in a search head cluster, in accordance with one or more aspects of the present disclosure. Referring to FIG. **6**, at block **610**, the search head leader monitors search head followers in the search head cluster. In one implementation, the search head leader sends out periodic heartbeat messages to maintain its authority. For example, the heartbeat messages may be transmitted approximately every 150 milliseconds. The heartbeat messages serve to inform the other search heads that the search head leader is operational. Accordingly, heartbeat messages may, in some circumstances, not be sent to a search head follower if the search head leader is sending other messages to the search head follower as these other messages serve to inform the search head leader that the search head follower is operational. The search head followers acknowledge the received heartbeat messages from the search head leader, with their own heartbeat messages sent back to the search head leader.

[0073] If a period of time passes without the search head leader receiving a heartbeat message from a search head follower, referred to as a heartbeat timeout, the search head leader may begin a node failure process. Thus, at decision block **620**, the search head leader determines whether a heartbeat message has been received for a search head follower. If a heartbeat message has been received, the search head leader returns to decision block **620** to await another heartbeat message. If a heartbeat

message has not been received, the search head leader continues to decision block **630** where it determines whether a heartbeat timeout has occurred. The heartbeat timeout may be randomly selected by the search head leader between a first value and a second value. For example, the heartbeat timeout may be a configured time period established by an administrator of the data aggregation and analysis system, such as 60 seconds.

[0074] If the search head leader determines, at decision block **630**, that a heartbeat timeout has not occurred, the search head leader returns to decision block **620** and awaits a heartbeat message. If the search head leader determines that a heartbeat timeout has occurred, the search head leader continues to block **640**, where it begins a search head failure process for purposes of updating search replication of search results. At block **640**, all references (e.g., listing or identifier of the failed search head) to the failed search head are removed from a cluster search results registry maintained by the search head leader. In one implementation, the cluster search results registry is the same as cluster search results registry **330** described with respect to FIG. **3A**.

[0075] Subsequently, at block **650**, a fix-up process is initiated for all modified search results from block **640**. The fix-up process refers to the process for placing a search result in compliance with the replication policy (e.g., replication count) configured for the search head cluster. For example, a fix-up process may determine that a replication count associated with a search head is not in compliance with the default replication count, and thus may schedule the search result for replication at one or more search head (if more replicas should be provided per portions of method **400** described above with respect to FIG. **4**) or may schedule the search result for removal of excess copies (e.g., per method **800** described below with respect to FIG. **8**).

[0076] FIG. **7** depicts a flow diagram of an example method **700** for removing excess copies of search results in association with search replication in a search head cluster, in accordance with one or more aspects of the present disclosure. A number of search result replications in the search head cluster may exceed the default replication count, for example, when a search result is replicated due to a proxy request, as described further below with respect to FIG. **8**. Referring to FIG. **7**, at block **720**, the search head leader parses a cluster search results registry to determine a number of copies of each search result. The number of copies may reflect both the original search result and the replicated search result(s), which are all identified in the cluster search results registry. Then, at decision block **730**, it is determined whether the number of copies of a search result exceeds a default replication count configured for the search head cluster. If not, then method **700** returns to decision block **730** to examine remaining search results.

[0077] If the number of search results copies does exceed the default replication count at decision block **730**, then method **700** proceeds to block **740**, where a least-recently-used replication for the search result is identified. The search head leader may be informed of recency of use statistics by the search heads each time a copy is accessed at the search head. In some implementations, the search head leader may maintain a counter for each search result copy that is incremented each time the search head leader receives an indication that the particular search head copy was accessed. The search head leader may then select the search result replication with the lowest counter value for removal. In some implementations, the original search result is not considered for excess replication removal. In other implementations, the original search result may be considered for excess replication removal.

[0078] At block **750**, the reference to the identified least-recently-used search result copy is removed from the cluster search results registry. Then, at block **760**, the search head storing the removed replication is instructed to delete the replicated search result copy from its search results storage.

[0079] Referring back to FIG. **3A**, the search head leader **104A** may also include a proxying component **320**. Proxying component **320** is responsible for enabling proxying of search result requests within the search head cluster **104** by managing updates of search result states to the other search heads in the search head cluster and, optionally, performing search result replication based

on results of proxy requests. Although replicating search results provides for high-availability of the search results within the cluster, the replication may not involve replicating search results to all search heads in the cluster, as this can be inefficient in terms of storage and network communications. However, when search results are replicated to a subset (e.g., a smaller set) of the search heads in the cluster, it may result in a user device communicating, during a user session, with a search head (referred to herein as the “local search head”) that does not maintain a local copy of the search result (either the originally-generated search result or a replicated version of the search result). Proxying of search requests in implementations of the disclosure enables the local search head to intelligently direct the user device to another search head that stores the specific search result replication. Further details of the replication process performed by a search head (e.g., local search head) are described with respect to FIG. 8.

[0080] FIG. 8 depicts a flow diagram of an example method **800** for search result proxying utilizing replicated search results in a search head cluster, in accordance with one or more aspects of the present disclosure. In some implementations, a proxying component **320** of a search head may perform the proxying process described herein. Referring to FIG. 8, at block **810**, a search result request is received from a user device at a local search head during a user session. In one implementation, the user device may include a browser, which communicates the search result request to the search head. At block **820**, the local search head determines whether a local copy of the request search results is stored at the local search head. The search result may be stored at the local search head if the local search head originated the search result or received a replica of the search result as part of the search result replication describe above. In one implementation, the local search head may reference its local cluster search result registry or may look directly at a file system of the local search head in order to determine if the local search head has a local copy of the search result.

[0081] If a local copy of the search result is found at decision block **830**, then method **800** proceeds to block **840** to serve the local copy of the search result to the user device. If a local copy of the search result is not found at decision block **830**, the method **800** continues to block **850**. At block **850**, the local cluster search result registry is referenced to identify one or more search heads of the search result cluster that have a copy of the search result. As discussed above, the local cluster search results registry may be a local copy that tracks a state of the primary cluster search results registry maintained by the search head leader. In one implementation, a proxying component **320**, as described with respect to FIG. 3A, of the search head leader may periodically send to each search head follower a state of the cluster search results registry maintained at the search head leader. The proxying component **320** may send a delta (e.g., rolling update) of the cluster search results registry to the search head follower so that the entire cluster search results registry is not sent each time; instead, updates (e.g., which search results have been replicated to which search heads, which search results have been deleted, etc.) from the previous transmission are provided.

[0082] Responsive to identifying search heads having the search result from the local cluster search results registry, a candidate selection process is applied to the identified search heads at block **860**. The candidate selection process selects one of the identified search heads to proxy the search result request. In one implementation, the candidate selection is site-aware so that the search head that is closest to the local search head in terms of site (e.g., geographic placement) is biased for selection purposes. If there is more than one search head equally close (e.g., at the same site) to the local search head, then the local search head may randomly select between the equally close search heads. If there are no search head candidates close to the local search head in terms of site, then the local search head may randomly select from the candidate search heads.

[0083] At block **870**, a proxy request for the search result is sent to the selected search head. Proxying may include, but is not limited to, a HyperText Transport Protocol (HTTP) proxy request, any type of redirect request that is natively supported by HTTP, a REST API call, and so on. Once the selected search head receives the proxy request, the search result is obtained from the selected

search head and served back to the user device.

[0084] In one implementation, if the local search head has no knowledge of the search result in its local cluster search results registry (e.g., the local cluster search results registry may not be up-to-date) and/or the selected search head is not responsive, the local search head may request search result location information from the search head leader. In one implementation, the proxying component **320** of the search head leader may respond to the proxy information request. In other implementations, when a proxy request to a selected search head fails, the local search head may select and try another search result location determined from the local cluster search results registry of the local search head.

[0085] In some implementations, the local search head may optionally, at block **880**, inform the search head leader that a proxy request was made to the selected search head. In addition, the local search head may inform the search head leader that the search results should be replicated to the local search head because of the proxy request. In one implementation, the search head leader may asynchronously replicate the search result to the local search head in response to the proxy request notification. If the search result is replicated, the search head leader may then perform method **700** of FIG. 7, for example, to remove the excess copies of the search result from the search head cluster. Once the search result is replicated, the local search head then finds the search result locally and does not have to proxy to obtain the search result.

[0086] The systems and methods described herein above may be employed by various data processing systems, e.g., data aggregation and analysis systems. In certain implementations, the example data aggregation and analysis system may perform search queries on performance data that is stored as “events,” wherein each event comprises a collection of performance data and/or diagnostic information that is generated by a computer system and is correlated with a specific point in time. In various illustrative examples, the data processing system may be represented by the SPLUNK® ENTERPRISE system produced by Splunk Inc. of San Francisco, California, to store and process performance data.

1.1 Overview

[0087] Modern data centers often comprise thousands of host computer systems that operate collectively to service requests from even larger numbers of remote clients. During operation, these data centers generate significant volumes of performance data and diagnostic information that can be analyzed to quickly diagnose performance problems. In order to reduce the size of this performance data, the data is typically pre-processed prior to being stored based on anticipated data-analysis needs. For example, pre-specified data items can be extracted from the performance data and stored in a database to facilitate efficient retrieval and analysis at search time. However, the rest of the performance data is not saved and is essentially discarded during pre-processing. As storage capacity becomes progressively cheaper and more plentiful, there are fewer incentives to discard this performance data and many reasons to keep it.

[0088] This plentiful storage capacity is presently making it feasible to store massive quantities of minimally processed performance data at “ingestion time” for later retrieval and analysis at “search time.” Note that performing the analysis operations at search time provides greater flexibility because it enables an analyst to search all of the performance data, instead of searching pre-specified data items that were stored at ingestion time. This enables the analyst to investigate different aspects of the performance data instead of being confined to the pre-specified set of data items that were selected at ingestion time.

[0089] However, analyzing massive quantities of heterogeneous performance data at search time can be a challenging task. A data center may generate heterogeneous performance data from thousands of different components, which can collectively generate tremendous volumes of performance data that can be time-consuming to analyze. For example, this performance data can include data from system logs, network packet data, sensor data, and data generated by various applications. Also, the unstructured nature of much of this performance data can pose additional

challenges because of the difficulty of applying semantic meaning to unstructured data, and the difficulty of indexing and querying unstructured data using traditional database systems.

[0090] These challenges can be addressed by using an event-based system, such as the SPLUNK® ENTERPRISE system produced by Splunk Inc. of San Francisco, California, to store and process performance data. The SPLUNK® ENTERPRISE system is the leading platform for providing real-time operational intelligence that enables organizations to collect, index, and harness machine-generated data from various websites, applications, servers, networks, and mobile devices that power their businesses. The SPLUNK® ENTERPRISE system is particularly useful for analyzing unstructured performance data, which is commonly found in system log files. Although many of the techniques described herein are explained with reference to the SPLUNK® ENTERPRISE system, the techniques are also applicable to other types of data server systems.

[0091] In the SPLUNK® ENTERPRISE system, performance data is stored as “events,” wherein each event comprises a collection of performance data and/or diagnostic information that is generated by a computer system and is correlated with a specific point in time. Events can be derived from “time series data,” wherein time series data comprises a sequence of data points (e.g., performance measurements from a computer system) that are associated with successive points in time and are typically spaced at uniform time intervals. Events can also be derived from “structured” or “unstructured” data. Structured data has a predefined format, wherein specific data items with specific data formats reside at predefined locations in the data. For example, structured data can include data items stored in fields in a database table. In contrast, unstructured data does not have a predefined format. This means that unstructured data can comprise various data items having different data types that can reside at different locations. For example, when the data source is an operating system log, an event can include one or more lines from the operating system log containing raw data that includes different types of performance and diagnostic information associated with a specific point in time. Examples of data sources from which an event may be derived include, but are not limited to: web servers; application servers; databases; firewalls; routers; operating systems; and software applications that execute on computer systems, mobile devices, and sensors. The data generated by such data sources can be produced in various forms including, for example and without limitation, server log files, activity log files, configuration files, messages, network packet data, performance measurements and sensor measurements. An event typically includes a timestamp that may be derived from the raw data in the event, or may be determined through interpolation between temporally proximate events having known timestamps.

[0092] The SPLUNK® ENTERPRISE system also facilitates using a flexible schema to specify how to extract information from the event data, wherein the flexible schema may be developed and redefined as needed. Note that a flexible schema may be applied to event data “on the fly,” when it is needed (e.g., at search time), rather than at ingestion time of the data as in traditional database systems. Because the schema is not applied to event data until it is needed (e.g., at search time), it is referred to as a “late-binding schema.”

[0093] During operation, the SPLUNK® ENTERPRISE system starts with raw data, which can include unstructured data, machine data, performance measurements or other time-series data, such as data obtained from weblogs, syslogs, or sensor readings. It divides this raw data into “portions,” and optionally transforms the data to produce timestamped events. The system stores the timestamped events in a data store, and enables a user to run queries against the data store to retrieve events that meet specified criteria, such as containing certain keywords or having specific values in defined fields. Note that the term “field” refers to a location in the event data containing a value for a specific data item.

[0094] As noted above, the SPLUNK® ENTERPRISE system facilitates using a late-binding schema while performing queries on events. A late-binding schema specifies “extraction rules” that are applied to data in the events to extract values for specific fields. More specifically, the extraction rules for a field can include one or more instructions that specify how to extract a value

for the field from the event data. An extraction rule can generally include any type of instruction for extracting values from data in events. In some cases, an extraction rule comprises a regular expression, in which case the rule is referred to as a “regex rule.”

[0095] In contrast to a conventional schema for a database system, a late-binding schema is not defined at data ingestion time. Instead, the late-binding schema can be developed on an ongoing basis until the time a query is actually executed. This means that extraction rules for the fields in a query may be provided in the query itself, or may be located during execution of the query. Hence, as an analyst learns more about the data in the events, the analyst can continue to refine the late-binding schema by adding new fields, deleting fields, or changing the field extraction rules until the next time the schema is used by a query. Because the SPLUNK® ENTERPRISE system maintains the underlying raw data and provides a late-binding schema for searching the raw data, it enables an analyst to investigate questions that arise as the analyst learns more about the events.

[0096] In the SPLUNK® ENTERPRISE system, a field extractor may be configured to automatically generate extraction rules for certain fields in the events when the events are being created, indexed, or stored, or possibly at a later time. Alternatively, a user may manually define extraction rules for fields using a variety of techniques.

[0097] Also, a number of “default fields” that specify metadata about the events rather than data in the events themselves can be created automatically. For example, such default fields can specify: a timestamp for the event data; a host from which the event data originated; a source of the event data; and a source type for the event data. These default fields may be determined automatically when the events are created, indexed or stored.

[0098] In some embodiments, a common field name may be used to reference two or more fields containing equivalent data items, even though the fields may be associated with different types of events that possibly have different data formats and different extraction rules. By enabling a common field name to be used to identify equivalent fields from different types of events generated by different data sources, the system facilitates use of a “common information model” (CIM) across the different data sources.

1.2 Data Server System

[0099] FIG. 9 presents a block diagram of an exemplary event-processing system **100**, similar to the SPLUNK® ENTERPRISE system. System **100** includes one or more forwarders **101** that collect data obtained from a variety of different data sources **105**, and one or more indexers **102** that store, process, and/or perform operations on this data, wherein each indexer operates on data contained in a specific data store **103**. These forwarders and indexers can comprise separate computer systems in a data center, or may alternatively comprise separate processes executing on various computer systems in a data center.

[0100] During operation, the forwarders **101** identify which indexers **102** will receive the collected data and then forward the data to the identified indexers. Forwarders **101** can also perform operations to strip out extraneous data and detect timestamps in the data. The forwarders next determine which indexers **102** will receive each data item and then forward the data items to the determined indexers **102**.

[0101] Note that distributing data across different indexers facilitates parallel processing. This parallel processing can take place at data ingestion time, because multiple indexers can process the incoming data in parallel. The parallel processing can also take place at search time, because multiple indexers can search through the data in parallel.

[0102] System **100** and the processes described below with respect to FIGS. 1-5 are further described in “Exploring Splunk Search Processing Language (SPL) Primer and Cookbook” by David Carasso, CITO Research, 2012, and in “Optimizing Data Analysis With a Semi-Structured Time Series Database” by Ledion Bitincka, Archana Ganapathi, Stephen Sorkin, and Steve Zhang, SLAML, 2010, each of which is hereby incorporated herein by reference in its entirety for all purposes.

1.3 Data Ingestion

[0103] FIG. **10** presents a flowchart illustrating how an indexer processes, indexes, and stores data received from forwarders in accordance with the disclosed embodiments. At block **201**, the indexer receives the data from the forwarder. Next, at block **202**, the indexer apportions the data into events. Note that the data can include lines of text that are separated by carriage returns or line breaks and an event may include one or more of these lines. During the apportioning process, the indexer can use heuristic rules to automatically determine the boundaries of the events, which for example coincide with line boundaries. These heuristic rules may be determined based on the source of the data, wherein the indexer can be explicitly informed about the source of the data or can infer the source of the data by examining the data. These heuristic rules can include regular expression-based rules or delimiter-based rules for determining event boundaries, wherein the event boundaries may be indicated by predefined characters or character strings. These predefined characters may include punctuation marks or other special characters including, for example, carriage returns, tabs, spaces or line breaks. In some cases, a user can fine-tune or configure the rules that the indexers use to determine event boundaries in order to adapt the rules to the user's specific requirements.

[0104] Next, the indexer determines a timestamp for each event at block **203**. As mentioned above, these timestamps can be determined by extracting the time directly from data in the event, or by interpolating the time based on timestamps from temporally proximate events. In some cases, a timestamp can be determined based on the time the data was received or generated. The indexer subsequently associates the determined timestamp with each event at block **204**, for example by storing the timestamp as metadata for each event.

[0105] Then, the system can apply transformations to data to be included in events at block **205**. For log data, such transformations can include removing a portion of an event (e.g., a portion used to define event boundaries, extraneous text, characters, etc.) or removing redundant portions of an event. Note that a user can specify portions to be removed using a regular expression or any other possible technique.

[0106] Next, a keyword index can optionally be generated to facilitate fast keyword searching for events. To build a keyword index, the indexer first identifies a set of keywords in block **206**. Then, at block **207** the indexer includes the identified keywords in an index, which associates each stored keyword with references to events containing that keyword (or to locations within events where that keyword is located). When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword.

[0107] In some embodiments, the keyword index may include entries for name-value pairs found in events, wherein a name-value pair can include a pair of keywords connected by a symbol, such as an equals sign or colon. In this way, events containing these name-value pairs can be quickly located. In some embodiments, fields can automatically be generated for some or all of the name-value pairs at the time of indexing. For example, if the string “dest=10.0.1.2” is found in an event, a field named “dest” may be created for the event, and assigned a value of “10.0.1.2.”

[0108] Finally, the indexer stores the events in a data store at block **208**, wherein a timestamp can be stored with each event to facilitate searching for events based on a time range. In some cases, the stored events are organized into a plurality of buckets, wherein each bucket stores events associated with a specific time range. This not only improves time-based searches, but it also allows events with recent timestamps that may have a higher likelihood of being accessed to be stored in faster memory to facilitate faster retrieval. For example, a bucket containing the most recent events can be stored as flash memory instead of on hard disk.

[0109] Each indexer **102** is responsible for storing and searching a subset of the events contained in a corresponding data store **103**. By distributing events among the indexers and data stores, the indexers can analyze events for a query in parallel, for example using map-reduce techniques, wherein each indexer returns partial responses for a subset of events to a search head that combines

the results to produce an answer for the query. By storing events in buckets for specific time ranges, an indexer may further optimize searching by looking only in buckets for time ranges that are relevant to a query.

[0110] Moreover, events and buckets can also be replicated across different indexers and data stores to facilitate high availability and disaster recovery as is described in U.S. patent application Ser. No. 14/266,812 filed on 30 Apr. 2014, and in U.S. application patent Ser. No. 14/266,817 also filed on 30 Apr. 2014.

1.4 Query Processing

[0111] FIG. **11** presents a flowchart illustrating how a search head and indexers perform a search query in accordance with the disclosed embodiments. At the start of this process, a search head receives a search query from a client at block **301**. Next, at block **302**, the search head analyzes the search query to determine what portions can be delegated to indexers and what portions need to be executed locally by the search head. At block **303**, the search head distributes the determined portions of the query to the indexers. Note that commands that operate on single events can be trivially delegated to the indexers, while commands that involve events from multiple indexers are harder to delegate.

[0112] Then, at block **304**, the indexers to which the query was distributed search their data stores for events that are responsive to the query. To determine which events are responsive to the query, the indexer searches for events that match the criteria specified in the query. This criteria can include matching keywords or specific values for certain fields. In a query that uses a late-binding schema, the searching operations in block **304** may involve using the late-binding scheme to extract values for specified fields from events at the time the query is processed. Next, the indexers can either send the relevant events back to the search head, or use the events to calculate a partial result, and send the partial result back to the search head.

[0113] Finally, at block **305**, the search head combines the partial results and/or events received from the indexers to produce a final result for the query. This final result can comprise different types of data depending upon what the query is asking for. For example, the final results can include a listing of matching events returned by the query, or some type of visualization of data from the returned events. In another example, the final result can include one or more calculated values derived from the matching events.

[0114] Moreover, the results generated by system **100** can be returned to a client using different techniques. For example, one technique streams results back to a client in real-time as they are identified. Another technique waits to report results to the client until a complete set of results is ready to return to the client. Yet another technique streams interim results back to the client in real-time until a complete set of results is ready, and then returns the complete set of results to the client. In another technique, certain results are stored as “search jobs,” and the client may subsequently retrieve the results by referencing the search jobs.

[0115] The search head can also perform various operations to make the search more efficient. For example, before the search head starts executing a query, the search head can determine a time range for the query and a set of common keywords that all matching events must include. Next, the search head can use these parameters to query the indexers to obtain a superset of the eventual results. Then, during a filtering stage, the search head can perform field-extraction operations on the superset to produce a reduced set of search results.

1.5 Field Extraction

[0116] FIG. **12** presents a block diagram illustrating how fields can be extracted during query processing in accordance with the disclosed embodiments. At the start of this process, a search query **402** is received at a query processor **404**. Query processor **404** includes various mechanisms for processing a query, wherein these mechanisms can reside in a search head **104** and/or an indexer **102**. Note that the exemplary search query **402** illustrated in FIG. **4** is expressed in Search Processing Language (SPL), which is used in conjunction with the SPLUNK® ENTERPRISE

system. SPL is a pipelined search language in which a set of inputs is operated on by a first command in a command line, and then a subsequent command following the pipe symbol “” operates on the results produced by the first command, and so on for additional commands. Search query **402** can also be expressed in other query languages, such as the Structured Query Language (“SQL”) or any suitable query language.

[0117] Upon receiving search query **402**, query processor **404** sees that search query **402** includes two fields “IP” and “target.” Query processor **404** also determines that the values for the “IP” and “target” fields have not already been extracted from events in data store **414**, and consequently determines that query processor **404** needs to use extraction rules to extract values for the fields. Hence, query processor **404** performs a lookup for the extraction rules in a rule base **406**, wherein rule base **406** maps field names to corresponding extraction rules and obtains extraction rules **408-409**, wherein extraction rule **408** specifies how to extract a value for the “IP” field from an event, and extraction rule **409** specifies how to extract a value for the “target” field from an event. As is illustrated in FIG. **12**, extraction rules **408-409** can comprise regular expressions that specify how to extract values for the relevant fields. Such regular-expression-based extraction rules are also referred to as “regex rules.” In addition to specifying how to extract field values, the extraction rules may also include instructions for deriving a field value by performing a function on a character string or value retrieved by the extraction rule. For example, a transformation rule may truncate a character string, or convert the character string into a different data format. In some cases, the query itself can specify one or more extraction rules.

[0118] Next, query processor **404** sends extraction rules **408-409** to a field extractor **412**, which applies extraction rules **408-409** to events **416-418** in a data store **414**. Note that data store **414** can include one or more data stores, and extraction rules **408-409** can be applied to large numbers of events in data store **414**, and are not meant to be limited to the three events **416-418** illustrated in FIG. **12**. Moreover, the query processor **404** can instruct field extractor **412** to apply the extraction rules to all the events in a data store **414**, or to a subset of the events that have been filtered based on some criteria.

[0119] Next, field extractor **412** applies extraction rule **408** for the first command “Search IP=“10*” to events in data store **414** including events **416-418**. Extraction rule **408** is used to extract values for the IP address field from events in data store **414** by looking for a pattern of one or more digits, followed by a period, followed again by one or more digits, followed by another period, followed again by one or more digits, followed by another period, and followed again by one or more digits. Next, field extractor **412** returns field values **420** to query processor **404**, which uses the criterion IP=“10*” to look for IP addresses that start with “10”. Note that events **416** and **417** match this criterion, but event **418** does not, so the result set for the first command is events **416-417**.

[0120] Query processor **404** then sends events **416-417** to the next command “stats count target.” To process this command, query processor **404** causes field extractor **412** to apply extraction rule **409** to events **416-417**. Extraction rule **409** is used to extract values for the target field for events **416-417** by skipping the first four commas in events **416-417**, and then extracting all of the following characters until a comma or period is reached. Next, field extractor **412** returns field values **421** to query processor **404**, which executes the command “stats count target” to count the number of unique values contained in the target fields, which in this example produces the value “2” that is returned as a final result **422** for the query.

[0121] Note that query results can be returned to a client, a search head, or any other system component for further processing. In general, query results may include: a set of one or more events; a set of one or more values obtained from the events; a subset of the values; statistics calculated based on the values; a report containing the values; or a visualization, such as a graph or chart, generated from the values.

1.6 Exemplary Search Screen

[0122] FIG. 14A illustrates an exemplary search screen **600** in accordance with the disclosed embodiments. Search screen **600** includes a search bar **602** that accepts user input in the form of a search string. It also includes a time range picker **612** that enables the user to specify a time range for the search. For “historical searches” the user can select a specific time range, or alternatively a relative time range, such as “today,” “yesterday” or “last week.” For “real-time searches,” the user can select the size of a preceding time window to search for real-time events. Search screen **600** also initially displays a “data summary” dialog as is illustrated in FIG. 14B that enables the user to select different sources for the event data, for example by selecting specific hosts and log files. [0123] After the search is executed, the search screen **600** can display the results through search results tabs **604**, wherein search results tabs **604** includes: an “events tab” that displays various information about events returned by the search; a “statistics tab” that displays statistics about the search results; and a “visualization tab” that displays various visualizations of the search results. The events tab illustrated in FIG. 14A displays a timeline graph **605** that graphically illustrates the number of events that occurred in one-hour intervals over the selected time range. It also displays an events list **608** that enables a user to view the raw data in each of the returned events. It additionally displays a fields sidebar **606** that includes statistics about occurrences of specific fields in the returned events, including “selected fields” that are pre-selected by the user, and “interesting fields” that are automatically selected by the system based on pre-specified criteria.

1.7 Acceleration Techniques

[0124] The above-described system provides significant flexibility by enabling a user to analyze massive quantities of minimally processed performance data “on the fly” at search time instead of storing pre-specified portions of the performance data in a database at ingestion time. This flexibility enables a user to see correlations in the performance data and perform subsequent queries to examine interesting aspects of the performance data that may not have been apparent at ingestion time.

[0125] However, performing extraction and analysis operations at search time can involve a large amount of data and require a large number of computational operations, which can cause considerable delays while processing the queries. Fortunately, a number of acceleration techniques have been developed to speed up analysis operations performed at search time. These techniques include: (1) performing search operations in parallel by formulating a search as a map-reduce computation; (2) using a keyword index; (3) using a high performance analytics store; and (4) accelerating the process of generating reports. These techniques are described in more detail below.

1.7.1 Map-Reduce Technique

[0126] To facilitate faster query processing, a query can be structured as a map-reduce computation, wherein the “map” operations are delegated to the indexers, while the corresponding “reduce” operations are performed locally at the search head. For example, FIG. 13 illustrates how a search query **501** received from a client at search head **104** can split into two phases, including: (1) a “map phase” comprising subtasks **502** (e.g., data retrieval or simple filtering) that may be performed in parallel and are “mapped” to indexers **102** for execution, and (2) a “reduce phase” comprising a merging operation **503** to be executed by the search head when the results are ultimately collected from the indexers.

[0127] During operation, upon receiving search query **501**, search head **104** modifies search query **501** by substituting “stats” with “prestats” to produce search query **502**, and then distributes search query **502** to one or more distributed indexers, which are also referred to as “search peers.” Note that search queries may generally specify search criteria or operations to be performed on events that meet the search criteria. Search queries may also specify field names, as well as search criteria for the values in the fields or operations to be performed on the values in the fields. Moreover, the search head may distribute the full search query to the search peers as is illustrated in FIG. 11, or may alternatively distribute a modified version (e.g., a more restricted version) of the search query to the search peers. In this example, the indexers are responsible for producing the results and

sending them to the search head. After the indexers return the results to the search head, the search head performs the merging operations **503** on the results. Note that by executing the computation in this way, the system effectively distributes the computational operations while minimizing data transfers.

1.7.2 Keyword Index

[0128] As described above with reference to the flow charts in FIGS. **10** and **11**, event-processing system **100** can construct and maintain one or more keyword indices to facilitate rapidly identifying events containing specific keywords. This can greatly speed up the processing of queries involving specific keywords. As mentioned above, to build a keyword index, an indexer first identifies a set of keywords. Then, the indexer includes the identified keywords in an index, which associates each stored keyword with references to events containing that keyword, or to locations within events where that keyword is located. When an indexer subsequently receives a keyword-based query, the indexer can access the keyword index to quickly identify events containing the keyword.

1.7.3 High Performance Analytics Store

[0129] To speed up certain types of queries, some embodiments of system **100** make use of a high performance analytics store, which is referred to as a “summarization table,” that contains entries for specific field-value pairs. Each of these entries keeps track of instances of a specific value in a specific field in the event data and includes references to events containing the specific value in the specific field. For example, an exemplary entry in a summarization table can keep track of occurrences of the value “94107” in a “ZIP code” field of a set of events, wherein the entry includes references to all of the events that contain the value “94107” in the ZIP code field. This enables the system to quickly process queries that seek to determine how many events have a particular value for a particular field, because the system can examine the entry in the summarization table to count instances of the specific value in the field without having to go through the individual events or do extractions at search time. Also, if the system needs to process all events that have a specific field-value combination, the system can use the references in the summarization table entry to directly access the events to extract further information without having to search all of the events to find the specific field-value combination at search time.

[0130] In some embodiments, the system maintains a separate summarization table for each of the above-described time-specific buckets that stores events for a specific time range, wherein a bucket-specific summarization table includes entries for specific field-value combinations that occur in events in the specific bucket. Alternatively, the system can maintain a separate summarization table for each indexer, wherein the indexer-specific summarization table only includes entries for the events in a data store that is managed by the specific indexer.

[0131] The summarization table can be populated by running a “collection query” that scans a set of events to find instances of a specific field-value combination, or alternatively instances of all field-value combinations for a specific field. A collection query can be initiated by a user, or can be scheduled to occur automatically at specific time intervals. A collection query can also be automatically launched in response to a query that asks for a specific field-value combination.

[0132] In some cases, the summarization tables may not cover all of the events that are relevant to a query. In this case, the system can use the summarization tables to obtain partial results for the events that are covered by summarization tables, but may also have to search through other events that are not covered by the summarization tables to produce additional results. These additional results can then be combined with the partial results to produce a final set of results for the query. This summarization table and associated techniques are described in more detail in U.S. Pat. No. 8,682,925, issued on Mar. 25, 2014.

1.7.4 Accelerating Report Generation

[0133] In some embodiments, a data server system such as the SPLUNK® ENTERPRISE system can accelerate the process of periodically generating updated reports based on query results. To accelerate this process, a summarization engine automatically examines the query to determine

whether generation of updated reports can be accelerated by creating intermediate summaries. (This is possible if results from preceding time periods can be computed separately and combined to generate an updated report. In some cases, it is not possible to combine such incremental results, for example where a value in the report depends on relationships between events from different time periods.) If reports can be accelerated, the summarization engine periodically generates a summary covering data obtained during a latest non-overlapping time period. For example, where the query seeks events meeting a specified criteria, a summary for the time period includes only events within the time period that meet the specified criteria. Similarly, if the query seeks statistics calculated from the events, such as the number of events that match the specified criteria, then the summary for the time period includes the number of events in the period that match the specified criteria.

[0134] In parallel with the creation of the summaries, the summarization engine schedules the periodic updating of the report associated with the query. During each scheduled report update, the query engine determines whether intermediate summaries have been generated covering portions of the time period covered by the report update. If so, then the report is generated based on the information contained in the summaries. Also, if additional event data has been received and has not yet been summarized, and is required to generate the complete report, the query can be run on this additional event data. Then, the results returned by this query on the additional event data, along with the partial results obtained from the intermediate summaries, can be combined to generate the updated report. This process is repeated each time the report is updated. Alternatively, if the system stores events in buckets covering specific time ranges, then the summaries can be generated on a bucket-by-bucket basis. Note that producing intermediate summaries can save the work involved in re-running the query for previous time periods, so only the newer event data needs to be processed while generating an updated report. These report acceleration techniques are described in more detail in U.S. Pat. No. 8,589,403, issued on Nov. 19, 2013, and U.S. Pat. No. 8,412,696, issued on Apr. 2, 2011.

1.8 Security Features

[0135] The SPLUNK® ENTERPRISE platform provides various schemas, dashboards and visualizations that make it easy for developers to create applications to provide additional capabilities. One such application is the SPLUNK® APP FOR ENTERPRISE SECURITY, which performs monitoring and alerting operations and includes analytics to facilitate identifying both known and unknown security threats based on large volumes of data stored by the SPLUNK® ENTERPRISE system. This differs significantly from conventional Security Information and Event Management (SIEM) systems that lack the infrastructure to effectively store and analyze large volumes of security-related event data. Traditional SIEM systems typically use fixed schemas to extract data from pre-defined security-related fields at data ingestion time, wherein the extracted data is typically stored in a relational database. This data extraction process (and associated reduction in data size) that occurs at data ingestion time inevitably hampers future incident investigations, when all of the original data may be needed to determine the root cause of a security issue, or to detect the tiny fingerprints of an impending security threat.

[0136] In contrast, the SPLUNK® APP FOR ENTERPRISE SECURITY system stores large volumes of minimally processed security-related data at ingestion time for later retrieval and analysis at search time when a live security threat is being investigated. To facilitate this data retrieval process, the SPLUNK® APP FOR ENTERPRISE SECURITY provides pre-specified schemas for extracting relevant values from the different types of security-related event data, and also enables a user to define such schemas.

[0137] The SPLUNK® APP FOR ENTERPRISE SECURITY can process many types of security-related information. In general, this security-related information can include any information that can be used to identify security threats. For example, the security-related information can include network-related information, such as IP addresses, domain names, asset identifiers, network traffic

volume, uniform resource locator strings, and source addresses. (The process of detecting security threats for network-related information is further described in U.S. patent application Ser. No. 13/956,252, and Ser. No. 13/956,262.) Security-related information can also include endpoint information, such as malware infection data and system configuration information, as well as access control information, such as login/logout information and access failure notifications. The security-related information can originate from various sources within a data center, such as hosts, virtual machines, storage devices and sensors. The security-related information can also originate from various sources in a network, such as routers, switches, email servers, proxy servers, gateways, firewalls and intrusion-detection systems.

[0138] During operation, the SPLUNK® APP FOR ENTERPRISE SECURITY facilitates detecting so-called “notable events” that are likely to indicate a security threat. These notable events can be detected in a number of ways: (1) an analyst can notice a correlation in the data and can manually identify a corresponding group of one or more events as “notable;” or (2) an analyst can define a “correlation search” specifying criteria for a notable event, and every time one or more events satisfy the criteria, the application can indicate that the one or more events are notable. An analyst can alternatively select a pre-defined correlation search provided by the application. Note that correlation searches can be run continuously or at regular intervals (e.g., every hour) to search for notable events. Upon detection, notable events can be stored in a dedicated “notable events index,” which can be subsequently accessed to generate various visualizations containing security-related information. Also, alerts can be generated to notify system operators when important notable events are discovered.

[0139] The SPLUNK® APP FOR ENTERPRISE SECURITY provides various visualizations to aid in discovering security threats, such as a “key indicators view” that enables a user to view security metrics of interest, such as counts of different types of notable events. For example, FIG. 15A illustrates an exemplary key indicators view **700** that comprises a dashboard, which can display a value **701**, for various security-related metrics, such as malware infections **702**. It can also display a change in a metric value **703**, which indicates that the number of malware infections increased by 63 during the preceding interval. Key indicators view **700** additionally displays a histogram panel **704** that displays a histogram of notable events organized by urgency values, and a histogram of notable events organized by time intervals. This key indicators view is described in further detail in pending U.S. patent application Ser. No. 13/956,338 filed Jul. 31, 2013.

[0140] These visualizations can also include an “incident review dashboard” that enables a user to view and act on “notable events.” These notable events can include: (1) a single event of high importance, such as any activity from a known web attacker; or (2) multiple events that collectively warrant review, such as a large number of authentication failures on a host followed by a successful authentication. For example, FIG. 15B illustrates an exemplary incident review dashboard **710** that includes a set of incident attribute fields **711** that, for example, enables a user to specify a time range field **712** for the displayed events. It also includes a timeline **713** that graphically illustrates the number of incidents that occurred in one-hour time intervals over the selected time range. It additionally displays an events list **714** that enables a user to view a list of all of the notable events that match the criteria in the incident attributes fields **711**. To facilitate identifying patterns among the notable events, each notable event can be associated with an urgency value (e.g., low, medium, high, critical), which is indicated in the incident review dashboard. The urgency value for a detected event can be determined based on the severity of the event and the priority of the system component associated with the event. The incident review dashboard is described further in “<http://docs.splunk.com/Documentation/PCI/2.1.1/User/IncidentReviewdashboard>.”

1.9 Data Center Monitoring

[0141] As mentioned above, the SPLUNK® ENTERPRISE platform provides various features that make it easy for developers to create various applications. One such application is the SPLUNK® APP FOR VMWARE®, which performs monitoring operations and includes analytics to facilitate

diagnosing the root cause of performance problems in a data center based on large volumes of data stored by the SPLUNK® ENTERPRISE system.

[0142] This differs from conventional data-center-monitoring systems that lack the infrastructure to effectively store and analyze large volumes of performance information and log data obtained from the data center. In conventional data-center-monitoring systems, this performance data is typically pre-processed prior to being stored, for example by extracting pre-specified data items from the performance data and storing them in a database to facilitate subsequent retrieval and analysis at search time. However, the rest of the performance data is not saved and is essentially discarded during pre-processing. In contrast, the SPLUNK® APP FOR VMWARE® stores large volumes of minimally processed performance information and log data at ingestion time for later retrieval and analysis at search time when a live performance issue is being investigated.

[0143] The SPLUNK® APP FOR VMWARE® can process many types of performance-related information. In general, this performance-related information can include any type of performance-related data and log data produced by virtual machines and host computer systems in a data center. In addition to data obtained from various log files, this performance-related information can include values for performance metrics obtained through an application programming interface (API) provided as part of the vSphere Hypervisor™ system distributed by VMware, Inc. of Palo Alto, California. For example, these performance metrics can include: (1) CPU-related performance metrics; (2) disk-related performance metrics; (3) memory-related performance metrics; (4) network-related performance metrics; (5) energy-usage statistics; (6) data-traffic-related performance metrics; (7) overall system availability performance metrics; (8) cluster-related performance metrics; and (9) virtual machine performance statistics. For more details about such performance metrics, please see U.S. Pat. No. 14/167,316 filed 29 Jan. 2014, which is hereby incorporated herein by reference. Also, see “vSphere Monitoring and Performance,” Update 1, vSphere 5.5, EN-001357-00, <http://pubs.vmware.com/vsphere-55/topic/com.vmware.ICbase/PDF/vsphere-esxi-vcenter-server-551-monitoring-performance-guide.pdf>.

[0144] To facilitate retrieving information of interest from performance data and log files, the SPLUNK® APP FOR VMWARE® provides pre-specified schemas for extracting relevant values from different types of performance-related event data, and also enables a user to define such schemas.

[0145] The SPLUNK® APP FOR VMWARE® additionally provides various visualizations to facilitate detecting and diagnosing the root cause of performance problems. For example, one such visualization is a “proactive monitoring tree” that enables a user to easily view and understand relationships among various factors that affect the performance of a hierarchically structured computing system. This proactive monitoring tree enables a user to easily navigate the hierarchy by selectively expanding nodes representing various entities (e.g., virtual centers or computing clusters) to view performance information for lower-level nodes associated with lower-level entities (e.g., virtual machines or host systems). Exemplary node-expansion operations are illustrated in FIG. 15C, wherein nodes 733 and 734 are selectively expanded. Note that nodes 731-739 can be displayed using different patterns or colors to represent different performance states, such as a critical state, a warning state, a normal state or an unknown/offline state. The case of navigation provided by selective expansion in combination with the associated performance-state information enables a user to quickly diagnose the root cause of a performance problem. The proactive monitoring tree is described in further detail in U.S. patent application Ser. No. 14/235,490 filed on 15 Apr. 2014, which is hereby incorporated herein by reference for all possible purposes.

[0146] The SPLUNK® APP FOR VMWARE® also provides a user interface that enables a user to select a specific time range and then view heterogeneous data, comprising events, log data and associated performance metrics, for the selected time range. For example, the screen illustrated in

FIG. 15D displays a listing of recent “tasks and events” and a listing of recent “log entries” for a selected time range above a performance-metric graph for “average CPU core utilization” for the selected time range. Note that a user is able to operate pull-down menus **742** to selectively display different performance metric graphs for the selected time range. This enables the user to correlate trends in the performance-metric graph with corresponding event and log data to quickly determine the root cause of a performance problem. This user interface is described in more detail in U.S. patent application Ser. No. 14/167,316 filed on 29 Jan. 2014, which is hereby incorporated herein by reference for all possible purposes.

[0147] FIG. **16** illustrates a diagrammatic representation of a computing device **1000** within which a set of instructions for causing the computing device to perform the methods discussed herein may be executed. The computing device **1000** may be connected to other computing devices in a LAN, an intranet, an extranet, and/or the Internet. The computing device **1000** may operate in the capacity of a server machine in client-server network environment. The computing device **1000** may be provided by a personal computer (PC), a set-top box (STB), a server, a network router, switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single computing device is illustrated, the term “computing device” shall also be taken to include any collection of computing devices that individually or jointly execute a set (or multiple sets) of instructions to perform the methods discussed herein. In illustrative examples, the computing device **1000** may implement the above described methods **400-800** for replicating search results in a search head cluster.

[0148] The example computing device **1000** may include a processing device (e.g., a general purpose processor) **1002**, a main memory **1004** (e.g., synchronous dynamic random access memory (DRAM), read-only memory (ROM)), a static memory **1006** (e.g., flash memory and a data storage device **1018**), which may communicate with each other via a bus **1030**.

[0149] The processing device **1002** may be provided by one or more general-purpose processing devices such as a microprocessor, central processing unit, or the like. In an illustrative example, the processing device **1002** may comprise a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets or processors implementing a combination of instruction sets. The processing device **1002** may also comprise one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), a network processor, or the like. The processing device **1002** may be configured to execute the methods **400-800** for replicating search results in a search head cluster, in accordance with one or more aspects of the present disclosure.

[0150] The computing device **1000** may further include a network interface device **1008**, which may communicate with a network **1020**. The computing device **1000** also may include a video display unit **1010** (e.g., a liquid crystal display (LCD) or a cathode ray tube (CRT)), an alphanumeric input device **1012** (e.g., a keyboard), a cursor control device **1014** (e.g., a mouse) and an acoustic signal generation device **1016** (e.g., a speaker). In one embodiment, video display unit **1010**, alphanumeric input device **1012**, and cursor control device **1014** may be combined into a single component or device (e.g., an LCD touch screen).

[0151] The data storage device **1018** may include a computer-readable storage medium **1028** on which may be stored one or more sets of instructions (e.g., instructions of the methods **400-800** for replicating search results in a search head cluster, in accordance with one or more aspects of the present disclosure) implementing any one or more of the methods or functions described herein. Instructions implementing methods **400-800** may also reside, completely or at least partially, within main memory **1004** and/or within processing device **1002** during execution thereof by computing device **1000**, main memory **1004** and processing device **1002** also constituting computer-readable media. The instructions may further be transmitted or received over a network **1020** via network

interface device **1008**.

[0152] While computer-readable storage medium **1028** is shown in an illustrative example to be a single medium, the term “computer-readable storage medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database and/or associated caches and servers) that store one or more sets of instructions. The term “computer-readable storage medium” shall also be taken to include any medium that is capable of storing, encoding or carrying a set of instructions for execution by the machine and that cause the machine to perform the methods described herein. The term “computer-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media and magnetic media.

[0153] Unless specifically stated otherwise, terms such as “updating,” “identifying,” “determining,” “sending,” “assigning,” or the like refer to actions and processes performed or implemented by computing devices that manipulate and transform data represented as physical (electronic) quantities within the computing device's registers and memories into other data similarly represented as physical quantities within the computing device memories or registers or other such information storage, transmission or display devices. Also, the terms “first,” “second,” “third,” “fourth,” etc. as used herein are meant as labels to distinguish among different elements and may not necessarily have an ordinal meaning according to their numerical designation.

[0154] Examples described herein also relate to an apparatus for performing the methods described herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general purpose computing device selectively programmed by a computer program stored in the computing device. Such a computer program may be stored in a computer-readable non-transitory storage medium.

[0155] The methods and illustrative examples described herein are not inherently related to any particular computer or other apparatus. Various general purpose systems may be used in accordance with the teachings described herein, or it may prove convenient to construct more specialized apparatus to perform the required method operations. The required structure for a variety of these systems will appear as set forth in the description above.

[0156] The above description is intended to be illustrative, and not restrictive. Although the present disclosure has been described with references to specific illustrative examples, it will be recognized that the present disclosure is not limited to the examples described. The scope of the disclosure should be determined with reference to the following claims, along with the full scope of equivalents to which the claims are entitled.

Claims

1. (canceled)

2. A method comprising: determining, by a search head leader of a search head cluster comprising a plurality of search heads, a failed search head of the plurality of search heads; based on determining the failed search head, updating a cluster search results registry maintained by the search head leader by removing references to the failed search head from the cluster search results registry; and in accordance with the updated cluster search results registry, managing compliance of a replication policy in association with a search result corresponding with the failed search head.

3. The method of claim 2, wherein the failed search head is determined based on an expiration of a period of time without the search head leader receiving a heartbeat message from the failed search head.

4. The method of claim 2, further comprising: communicating, from the search head leader, periodic heartbeat messages to the failed search head; and determining, by the search head leader, an expiration of a period of time during which the search head leader fails to receive a heartbeat message from the failed search head.

5. The method of claim 2, wherein the references to the failed search head comprise an identifier

associated with the failed search head.

6. The method of claim 2, wherein removing the references to the failed search head from the cluster search results registry modifies one or more search results included in the cluster search results registry.

7. The method of claim 2, wherein the replication policy comprises a replication count configured in association with the search head cluster.

8. The method of claim 2, wherein the retention policy indicates a number of locations in the search head cluster at which the search result is to be replicated.

9. The method of claim 2, wherein managing compliance of the replication policy in association with the search result corresponding with the failed search head comprises determining that a replication count associated with the search result is not in compliance with the replication policy.

10. The method of claim 2, wherein managing compliance of the replication policy in association with the search result corresponding with the failed search head comprises: determining that a replication count associated with the search result is below the replication policy; and based on the determination that the replication count is below the replication policy, scheduling a replication of the search result to at least one search head of the plurality of search heads.

11. The method of claim 2, wherein managing compliance of the replication policy in association with the search result corresponding with the failed search head comprises: determining that a replication count associated with the search result is above the replication policy; and based on the determination that the replication count is below the replication policy, scheduling a removal of the search result from at least one search head of the plurality of search heads.

12. A system comprising: a memory; and one or more processing devices coupled with the memory, the one or more processing devices configured to: determine, by a search head leader of a search head cluster comprising a plurality of search heads, a failed search head of the plurality of search heads; based on determining the failed search head, update a cluster search results registry maintained by the search head leader by removing references to the failed search head from the cluster search results registry; and in accordance with the updated cluster search results registry, manage compliance of a replication policy in association with a search result corresponding with the failed search head.

13. The system of claim 12, wherein the failed search head is determined based on an expiration of a period of time without the search head leader receiving a heartbeat message from the failed search head.

14. The system of claim 12, wherein the one or more processing devices are further configured to: communicate, from the search head leader, periodic heartbeat messages to the failed search head; and determine, by the search head leader, an expiration of a period of time during which the search head leader fails to receive a heartbeat message from the failed search head.

15. The system of claim 12, wherein the references to the failed search head comprise an identifier associated with the failed search head.

16. The system of claim 12, wherein removing the references to the failed search head from the cluster search results registry modifies one or more search results included in the cluster search results registry.

17. A non-transitory computer-readable storage medium encoding executable instructions thereon that, in response to execution by one or more processing devices, cause the one or more processing devices to: determine, by a search head leader of a search head cluster comprising a plurality of search heads, a failed search head of the plurality of search heads; based on determining the failed search head, update a cluster search results registry maintained by the search head leader by removing references to the failed search head from the cluster search results registry; and in accordance with the updated cluster search results registry, manage compliance of a replication policy in association with a search result corresponding with the failed search head.

18. The non-transitory computer-readable storage medium of claim 17, wherein the replication

policy comprises a replication count configured in association with the search head cluster.

19. The non-transitory computer-readable storage medium of claim 17, wherein the retention policy indicates a number of locations in the search head cluster at which the search result is to be replicated.

20. The non-transitory computer-readable storage medium of claim 17, wherein managing compliance of the replication policy in association with the search result corresponding with the failed search head comprises determining that a replication count associated with the search result is not in compliance with the replication policy.

21. The non-transitory computer-readable storage medium of claim 17, wherein managing compliance of the replication policy in association with the search result corresponding with the failed search head comprises: determining that a replication count associated with the search result is below the replication policy; and based on the determination that the replication count is below the replication policy, scheduling a replication of the search result to at least one search head of the plurality of search heads.
