

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265357

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Sadeh; Norman

IMPLEMENTING A PRIVACY INFRASTRUCTURE FOR THE INTERNET OF THINGS

Abstract

An Internet of Things (IoT) resource remotely senses data about a user. A computing device of the user comprises a personal privacy app (“PPA”) that: receives data about the IoT resource, wherein the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; communicates a privacy request for the user with respect to the IoT resource, wherein the privacy request is one of the available user-specific privacy requests and wherein the privacy request communicated by the PPA is based on the data received about the IoT resource; and in response to a query related to the privacy request, causes electronic documentation to be transmitted that demonstrates that the user qualifies to submit the privacy request, such that the privacy request is applied to data collected about the user by the IoT resource.

Inventors: Sadeh; Norman (Pittsburgh, PA)

Applicant: Carnegie Mellon University (Pittsburgh, PA)

Family ID: 1000008575290

Assignee: Carnegie Mellon University (Pittsburgh, PA)

Appl. No.: 19/063255

Filed: February 25, 2025

Related U.S. Application Data

parent US continuation 17165775 20210202 parent-grant-document US 11768949 child US 18239267

parent US continuation 15858261 20171229 parent-grant-document US 10956586 child US 17165775

parent US continuation-in-part 18239267 20230829 parent-grant-document US 12235973 child US 19063255

Publication Classification

Int. Cl.: **G06F21/60** (20130101); **G06F21/62** (20130101); **G06N7/00** (20230101); **G06N20/00** (20190101); **H04L9/40** (20220101); **H04L67/12** (20220101); **H04L67/306** (20220101); **H04L67/50** (20220101)

U.S. Cl.:

CPC **G06F21/604** (20130101); **G06F21/6245** (20130101); **G06F21/629** (20130101); **G06N7/00** (20130101); **G06N20/00** (20190101); **H04L63/101** (20130101); **H04L63/102** (20130101); **H04L67/306** (20130101); **H04L67/12** (20130101); **H04L67/535** (20220501)

Background/Summary

PRIORITY CLAIM [0001] The present application claims priority as a continuation-in-part to U.S. nonprovisional patent application Ser. No. 18/239,267, filed Aug. 29, 2023, titled “Personalized Privacy Assistant,” which is a continuation of U.S. nonprovisional patent application Ser. No. 17/165,775, filed Feb. 2, 2021, titled “Personalized Privacy Assistant,” now U.S. Pat. No. 11,768,949, issued Sep. 26, 2023, which is a continuation of U.S. nonprovisional patent application Ser. No. 15/858,261, filed Dec. 29, 2017, titled “Personal Privacy Assistant,” now U.S. Pat. No. 10,956,586, issued Mar. 23, 2021, which is a continuation-in-part to U.S. nonprovisional patent application Ser. No. 15/658,204, filed Jul. 24, 2017, titled “Personalized Privacy Assistant,” which claims priority to U.S. provisional patent application Ser. No. 62/493,972, filed Jul. 22, 2016, all of which are incorporated herein by reference in their entirety.

BACKGROUND

[0003] Mobile app ecosystems such as Android or iOS compete in part based on the number, and the quality of apps they offer. To attract developers and help generate more apps, these platforms have exposed a growing number of Application Programming Interfaces (APIs). These APIs provide access to smartphone functionality (e.g., GPS, accelerometer, camera) and user data (e.g., unique identifiers, location, social media accounts), much of which is privacy-sensitive.

[0004] The Android and iOS platforms both rely on permission-based mechanisms that allow users to control access to sensitive data and functionality. While providing users with this level of control is important, the end result is an unwieldy number of app-permission decisions that users are expected to make. Estimates indicate that users, on average, have to make over one hundred permission decisions (e.g., typical users may have 50 installed apps on their phones, with many apps requiring 2 or more permissions). Prior work has shown that users are often unaware of, if not uncomfortable with many of the permissions they have ostensibly consented to at some point.

[0005] This situation is not unique to mobile apps on smartphones. It reflects the increasingly diverse ways in which technologies can access sensitive data and functionality and the need for users to be able to control settings that determine which technology can access which element of functionality and which data under which conditions. Other examples include security and privacy settings found in browsers, privacy settings associated with social network sites and applications, as well as emerging privacy settings associated with Internet of Things Resources (“IoT resources”). Permissions such as those found in iOS and Android are a mechanism for providing users with such control. If all users felt the same way about which sensitive data or functionality

different technologies can access such settings would not be required. Instead access could be configured to align with everyone's privacy preferences, namely their comfort in granting access to sensitive functionality and data under different conditions. Experiments show however that people's privacy preferences are diverse and that "one-size-fits-all" settings would often not do a good job at capturing these preferences, thus the need for user-configurable settings. As already pointed out the number of such settings has been growing and can be expected to continue to grow over time, leading to a situation where the number of such settings is unrealistically large for people to manage.

SUMMARY

[0006] The present invention revolves around, in one general aspect, personalized privacy assistant functionality that can help users configure privacy settings. This functionality is intended to reduce user burden and increase the accuracy with which people's permission settings reflect their actual privacy preferences. This functionality optionally takes advantage of machine learning and/or statistical analysis techniques and can be extended to a number of technologies and environments (e.g. configuring people's privacy settings in browsers, in Internet of Things (IoT) environments, on social networking sites and more). It can also be further extended to notifying users about privacy practices they may not be expecting or with which they are believed to be particularly uncomfortable, even when settings to turn them off or opt out of them are not available.

[0007] In one general aspect, therefore, the present invention is directed to systems and methods that configure permission settings for applications ("apps") running on a computing device of a user. The system may comprise a data center that generates at least one model of privacy preferences based on data collected from a population of user. The computing device is in communication with the data center via a communications network. The computing device comprises a processor that executes at least a first app that requests access to at least one permission of the computing device. The processor also executes the personal privacy assistant app. The personal privacy assistant app, in one embodiment, receives and stores locally the model from the data center. It also collects information about the user and, based thereon, identifies at least one recommended permission setting for the first app using the model and such that the recommended permission setting is user-specific. Then the personal privacy assistant app can configure the computing device to implement the recommended permission setting.

[0008] In another embodiment, the personal privacy assistant collects the user information and transmits it to the data center, which identifies the recommended permission settings(s) and transmits them back the personal privacy app for configuration. In yet another embodiment, the personal privacy assistant collects user information and stores it locally, using stored preferences to determine what permissions settings to recommend to the user or directly configure these permission settings. Collecting user information may take place by relying on user manually providing this information. In some embodiments such collection can also involve dialogues with the user. In some embodiments, user information may also involve looking at decisions made by the user (e.g. privacy settings already selected by the user), observing the user's behavior. It may also involve inferring information about the user based on any of the above information.

[0009] The personal privacy assistant can also be used for configuring user-specific privacy settings for IoT resources running on third party computing systems. For example, the system could comprise one or more IoT resources and at least a first IoT resources registry that advertises the one or more IoT resources. The personal privacy assistant running on the user's computing device receives from the first IoT resource registry data about the one or more IoT resources advertised by the IoT resource registry and configures user-specific permission settings for the one or more IoT resources based on the data received from the IoT resource registry and based on individual privacy preferences for the user stored by the personal privacy assistant, whether locally or in a decentralized manner, including possibly in the cloud. The data about the one or more IoT resources received by the personal privacy assistant may comprise privacy opt-in or privacy opt-out

settings (or other privacy settings) available for the one or more IoT resources. The personal privacy assistant may communicate directly with the one or more IoT resources, with policy enforcement points associated with one or more IoT resources to configure the user-specific privacy settings for these one or more IoT resources. This includes communication with consent management platforms and other similar systems.

[0010] In one general aspect, the present invention is directed to a system: comprising an Internet of Things (IoT) resource for remotely sensing data about a user; and a computing device of the user. The computing device comprises a processor that executes a personal privacy app (“PPA”). The PPA is configured to: (i) receive data about the IoT resource, where the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; (ii) communicate a privacy request for the user with respect to the IoT resource, where the privacy request is one of the available user-specific privacy requests and where the privacy request communicated by the PPA is based on the data received about the IoT resource; and (iii) in response to a query related to the privacy request, causes electronic documentation to be transmitted that demonstrates that the user qualifies to submit the privacy request, such that the privacy request is applied to data collected about the user by the IoT resource.

[0011] In various implementations, the PPA may communicate the privacy request to an authorized agent, which can forward the privacy request for implementation with respect to the data collected about the user by the IoT resource. There can be one or more authorized agents available for a given IoT resource, and identification of those one or more authorized agents can be included in the IoT resource description. Also, the APIs to communicate with the one or more authorized agents can also be included in the IoT resource description. Thus, for instance, the PPA might discover an authorized agent for a given IoT resource, it might then contact that authorized agent to find what specific privacy requests are supported for that resource, it might then combine this information with the data practice information contained in the IoT resource description and integrate this information for display in the PPA to the user (or for the PPA to automatically communicate specific privacy requests on behalf of the user based on the user's privacy preferences).

[0012] In another general aspect, the inventive system comprises: the IoT resource for remotely sensing data about a user; a computing device of the user, where the computing device comprises a processor that executes a personal privacy app (“PPA”); and an authorized agent that is authorized to submit privacy requests on behalf of the user and is a trusted source of privacy requests for the IoT resource. The PPA is configured to receive data about the IoT resource, where the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; and communicate, to the authorized agent, a privacy request for the user with respect to the IoT resource, where the privacy request is based on the data received about the IoT resource. Also, the authorized agent is configured to communicate the privacy request for the user for implementation with respect to data collected about the user by the IoT resource.

[0013] Data about the IoT resource could come entirely from the IoT resource description or may also include additional data obtained indirectly by using APIs provided in the IoT resource description and using these APIs to contact the IoT resource and obtain additional information about the resource's data practices, such as the particular privacy requests supported by the resource, the documentation required to demonstrate that the user qualifies to submit a given privacy request, and/or the default practices (e.g., privacy settings) of the IoT resource.

[0014] In various implementations, in response to a query related to the privacy request, the PPA causes electronic documentation to be transmitted to the authorized agent that demonstrates that the user qualifies to submit the privacy request. The authorized agent can then forward the documentation, such that the privacy request is applied to data collected about the user by the IoT resource. In various implementations, the list of available privacy requests may be communicated to the PPA by the IoT privacy infrastructure, such as an IoT resource registry and/or the IoT resource description hosted by the IoT privacy infrastructure. In various implementations, the

electronic documentation that shows that the user is qualified to submit the privacy could comprise a digital image (e.g., of the user's driver's license, etc.) or an electronic document with a digital signature to verify authenticity of the document. In various embodiments, the PPA is configured to determine the privacy request using a privacy preference model for the user. In various implementations, the available user-specific privacy requests comprise opt-in and opt-out settings for data practices of the IoT resource, to request deletion of the data collected about the user, and/or to review (such as by being sent a copy of) the data collected about the user by the IoT resource. In various implementations, the PPA also receiving an electronic notification that the privacy setting was received, verified, and/or implemented with respect to the data collected about the user by the IoT resource. In various embodiments, the system further comprises an IoT resource registry that advertises the IoT resource to the PPA. In which case, the PPA can discover the IoT resource registry based on location of the user computing device.

[0015] Embodiments of the present invention can allow users to more accurately, efficiently and easily align the privacy-related permissions actually enforced by their computing devices or by third party computing systems such as IoT resources with their privacy preferences. These and other benefits of the present invention will be apparent from the description below.

Description

FIGURES

[0016] Various embodiments of the present invention are described herein by way of example in connection with the following figures.

[0017] FIGS. **1A** and **1B** are diagrams illustrating two different ways of deploying a personalized privacy assistant to help users configure privacy settings according to various embodiments of the present invention.

[0018] FIG. **2** is a block diagram of a system according to various embodiments of the present invention, including a user computing device that hosts the Personalized Privacy Assistant—and communicates with a remote data center.

[0019] FIG. **3** shows functions of the Personalized Privacy Assistant **230** according to various embodiments.

[0020] FIG. **4** is a screen shot showing an example of a privacy nudge provided by the personalized privacy assistant according to various embodiments of the present invention.

[0021] FIGS. **5**, **8A-8C** and **9** are screen shots of example interfaces through which the user can accept or deny the recommended permission settings according to various embodiments of the present invention.

[0022] FIG. **6** is a collection of grids that show the variations between sample privacy preference profiles associated to different clusters of users, according to various embodiments of the present invention, with users in different profiles feeling more or less comfortable granting permissions to apps in different categories.

[0023] FIGS. **7A-7C** are screen shots of example questions that can be provided by the personalized privacy assistant in order to collect information about the privacy preferences of the user according to various embodiments of the present invention. These figures correspond to an embodiment of the invention deployed in the form of a mobile app in the +Play Store.

[0024] FIG. **10** illustrates an example of how additional personalized questions can be asked of the user after the user accepts, rejects or modifies a recommended setting according to various embodiments of the present invention.

[0025] FIG. **11** illustrates example app performance data that can be used in making the recommended permission settings according to various embodiments of the present invention.

[0026] FIGS. **12** and **13** and flow charts illustrating example process flows for embodiments of the

present invention.

[0027] FIGS. **14**, **15**, **16** illustrate example permission settings. FIG. **14** shows example permission settings for two different clusters. FIG. **15** illustrates example permission settings for a cluster where the purpose of the permission is used for the recommendations. And FIG. **16** illustrates example permission settings for a cluster where the context of the permission is used for the recommendations.

[0028] FIG. **17** illustrates a system for communicating a user privacy request for an IoT resource according to various embodiments of the present invention.

[0029] FIGS. **18A-E** illustrate an exemplary online template for which an owner or administrator of an IoT resource to publicize information about the IoT resource, according to various embodiments of the present invention.

[0030] FIGS. **19A-19E** and FIG. **22A-F** illustrate exemplary screenshots provided by a portal hosted by the IoT privacy infrastructure of FIG. **17** according to various embodiments of the present invention.

[0031] FIGS. **20A-D** and **23** illustrate exemplary screenshots provided by a dashboard hosted by the IoT privacy infrastructure of FIG. **17** for a user of the PPA according to various embodiments of the present invention.

[0032] FIG. **21** is a screenshot of the PPA available in an app store according to various embodiments of the present invention.

[0033] FIGS. **24A** and **24B** illustrate a clustering analysis that can be used to organize users in groups of like-minded users according to various embodiments of the present invention.

DESCRIPTION

[0034] In one general aspect, the present invention is directed to a personalized privacy assistant (PPA) that helps a user configure permission settings associated with technologies with which the user interacts. In one embodiment these technologies are applications (“apps”) running on a computing device such as a mobile computing device, such as a laptop, a tablet, or a wearable computing device. In other embodiments these permission settings are associated with services with which the user interacts such as cloud-based services (e.g., social networking site) or other computing systems, applications or devices such as a location tracking system deployed in a building, a camera monitoring system deployed in a department store, or IoT devices, applications or systems deployed in one's home (e.g. Amazon echo or other home assistant technology), a robot, or a car fitted with a number of sensors and other subsystems.

[0035] The permission settings control access by these technologies (e.g. apps on a computing device, IoT service, IoT app, IoT device, etc.) to sensitive data or functionality. An example is a set of permissions that enable a mobile operating system to control access to sensitive functionality or data by apps running on a smartphone (e.g., access to location tracking functionality, access to camera functionality, access to microphone functionality, access to messaging functionality, access to contacts list information, access to a device's unique identifier, access to health information, etc.). In another example, permissions control privacy and/or security settings associated with a browser or control which sensitive functionality or data websites can access. In yet another example, permissions may be used to capture and enforce user-specific settings associated with Internet of Things (IoT) devices or services such as opt-in or opt-out privacy settings associated with location tracking functionality in an office building, or permission to apply facial recognition functionality or scene recognition functionality to footage captured by a video monitoring system in a mall.

[0036] FIGS. **1A** and **1B** are diagrams illustrating two different ways of deploying a personalized privacy assistant to help users configure privacy settings according to various embodiments of the present invention. FIG. **1A** illustrates a deployment of a Personalized Privacy Assistant **230** to configure permission settings **231** on a user computing device **114** such as a smartphone or tablet, with the permission settings controlling access to both sensitive on-device functionality **140** (e.g.

GPS, camera, microphone) and sensitive on-device data **141** (e.g. contacts list, health information, IMEI number), and sensitive off-device (or external) third party data **151** and functionality **150** (e.g., social media data and functionality (such as accessing the user's Facebook account and/or posting on the user's Facebook account), third party payment data and functionality, information from third party health and fitness devices, external calendar or email services for the user). FIG. **1B** illustrates a deployment of a Personalized Privacy Assistant **230** to help users configure user-specific privacy settings associated with different external IoT resources such as IoT apps, IoT devices or IoT services, according to various embodiments of the present invention.

[0037] As shown in FIGS. **1A** and **1B**, the Personalized Privacy Assistant (PPA) may be itself an application running on a computing device operated by the user (e.g. a desktop, a laptop, a tablet, a smartphone, a wearable computing device, a smart TV, a smart fridge, a smart car, a robot), or it may be running on a server (e.g. a server running in the cloud) and simply interact with the user via an application such as a browser. The PPA's functionality may also be distributed across any number of computing systems with some of its functionality running locally on a computing device in the user's possession and some of its functionality running on one or more servers.

[0038] The Personalized Privacy Assistant maintains models of user privacy preferences. In particular it maintains models of preferences that reflect people's comfort in granting different permissions to different technologies under different conditions. In one embodiment these are models of people's willingness to grant permissions controlling access to sensitive functionality or data by apps running on a user's computing device. These models are typically built and refined by collecting and analyzing privacy preferences for a population of users (or possibly a subset of people in that population, sometimes referred to herein as "test subjects"). The collection of these privacy preferences will typically (though not necessarily) require obtaining some form of consent from the "test subjects." These collective privacy preference models, which may for instance come in the form of clusters of like-minded users and profiles of typical preferences associated with people in a given cluster, can be used to quickly match individual users against them and identify a number of permission settings that are likely to be well aligned with that user's privacy preferences. When such profiles are identified for a given user, they can be used to recommend permission settings to the user and/or automatically configure some of the user's settings, possibly after checking with the user. The Personalized Privacy Assistant may also maintain models of preferences for when and how often users want to be notified about the presence of technologies collecting and/or using their data in different context (e.g. notification that a mobile app just accessed one's location, notification that one is about to enter an area under camera surveillance or a building with location tracking functionality). Again these models are typically built and refined by collecting and analyzing privacy preferences for a population of users (or possibly a subset of people in that population, sometimes referred to herein as "test subjects") and matching them against information provided by the user, or they may simply be specified directly by the user.

[0039] Different methods for identifying profiles (e.g., a profile associated with a cluster of like-minded users) that best match a user's preferences and for helping users configure their privacy preferences based on these profiles are disclosed herein and a person skilled in the art will easily appreciate that many variations of these methods can be used for this purpose. Other methods for maintaining and using collective privacy preference models are also disclosed that do not rely on clustering techniques. In general these models can be developed and refined through a number of different machine learning and/or statistical analysis techniques. In various embodiments privacy assistants also develop and refine individual privacy preference models, whether to replace or supplement collective privacy preference models. For instance, collective privacy preference models can be used to recommend an initial set of permission settings to users. In a later phase, information about the particular permission setting recommendations a user accepts, rejects, or modifies as well as additional permission settings a user may select, can be used to develop finer, individual privacy preference models, which in turn can be used to provide users with additional

recommendations over time.

[0040] A challenge in building collective privacy preferences for users is that many users are not aware of and/or do not bother to review and configure their permission settings. Accordingly, applying machine learning techniques and/or statistical analysis techniques directly to settings collected from a population of users (or a subset of users in that population) may produce privacy models that are not reflecting people's true privacy preferences. In some embodiments, the present invention relies on post-processing techniques, where settings collected from test subjects are only kept for analysis if data collected about a subject indicates that he or she was truly engaged with their settings (e.g., the user modified a threshold number of settings, or the user at least reviewed his or her settings a threshold number of times, possibly over a given period of time).

[0041] In some embodiments, the present invention also relies on “privacy nudges” intended to motivate users to review and configure their permission settings and increase the chance that the settings are well aligned with their privacy preferences. Once users have been subjected to a sufficient number of nudges, their permission settings can be collected and analyzed to develop stronger privacy models (both collective and individual privacy models), namely models that are likely to better capture people's privacy preferences when it comes to granting or denying different permissions. Privacy nudges can come in many different forms (e.g. pop-up message informing users of one or more facts that are likely to peak their interest/surprise them, or more generally are likely to motivate them to think more carefully about the ramifications of one or more privacy decisions associated with the configuration of some of their permissions). For instance, prior research has shown that mobile app privacy nudges can lead many users to (re) visit their mobile app permission settings and often modify a number of these settings, as they realize that they are not well aligned with their privacy preferences. Examples of information shown to have meaningful impact on users include information about the frequency at which some sensitive functionality or data has been accessed, the particular apps responsible for such access, the purpose for accessing this sensitive functionality or data, including who the data might be shared with, what information might be inferred from the sensitive data and what consequences this might have (e.g. loss of one's insurance policy if sensitive health information is shared with a health insurance company, speeding ticket if someone's driving speed is shared with the police, information about one's sexual orientation or religious affiliation being shared with parties a user might not feel comfortable sharing this information with, inferring one's income or education level, and much more).

[0042] In yet some other embodiments, privacy preferences from people can be collected by asking them to review a series of scenarios and asking them to indicate their level of comfort granting different permissions or privacy settings under these scenarios. Such an approach can lend itself to crowdsourcing privacy preference information online from a large number of subjects, possibly subject to compensation (e.g. using Amazon Turk or equivalent functionality).

[0043] FIG. 2 is a block diagram of one embodiment of a system in which the present invention can be employed. FIG. 2 illustrates a particular configuration of a Personalized Privacy Assistant **230** and a data center **106** (having one or more data servers) with which the Personalized Privacy Assistant **230** communicates to develop and refine models of people's privacy preferences and recommend privacy settings (e.g. mobile app permission settings). As shown in FIG. 2, a data center **106** is in communication with a computing device **114** via a communications network **112** (e.g., the Internet). The computing device **114** may run or execute the personalized privacy assistant (PPA) as shown in FIG. 1, as described further below. The computing device may also run or execute one or more apps whose permissions are controlled with the aid of the PPA.

[0044] A number of configurations are possible, with processing and storage being distributed differently between the user computing device **114** and the data center **106**. For instance, FIG. 1B illustrates a configuration in which the user computing device hosts the Personal Privacy Assistant (PPA) **230**. In other configurations, some or possibly even all of the Personal Privacy Assistant

functionality could be hosted in the data center **106**. The diagram in FIG. **2** illustrates how privacy preference modeling functionality **109** can be divided between the data center **106** and the computing device **114**. In this particular configuration, the data center **106** is responsible for building and refining collective privacy preference models **111** based on the collected user population privacy data **104**. In other embodiments, some or all of the functionality and models depicted in FIG. **2** as being hosted in the data center could actually reside on the user computing device **114**. Variations of this general system that are also covered by the present invention are described herein.

[0045] The data center **106** comprises one or more servers **108**. The data center **106** also comprises a database **104** that stores data about the privacy preferences of a population of users, which can be collected as described below. From this data, in one embodiment, privacy preference modeling techniques **109** (e.g., statistical and/or machine learning techniques, as further detailed below) can derive a collection of collective privacy preference profiles **111**. The collective privacy preference profiles can comprise collections of privacy preferences found to capture the way a significant majority of users in a given subset of the population of users (e.g., a cluster of like-minded users) feel about granting or denying certain permissions to different apps or categories of apps, and more generally to different technologies (e.g., to different IoT resources). These preference profiles could also capture preferences users in a subset of the population have when it comes to being notified about particular data collection or use practices in different contexts (e.g., users who want to be notified about the presence of cameras coupled with facial expression recognition technology, users who do not care to be notified about the presence of such cameras unless they are at a bar, or users who want to be notified when their mobile app shares their speed with the police). More details about the user privacy preferences for apps and other technologies are described below. Also, the servers **108** of the data center **106** could be co-located at one geographic center or distributed across multiple geographic locations, still collectively constituting the data center **106**. For example, one set of servers **108** could generate the models and the models could be made available to the PPAs on another, distinct set of servers **108**. Each set of servers **108** would collectively comprise the data center **106**.

[0046] The data center **106** can also comprise an app performance database **113** that stores data about the performance of apps, particularly the operational performance when permission settings are changed (e.g., whether the apps crash, lose some of their functionality, or impact other performance attributes, such as the battery life of the computing device **114** or the performance of other apps on the computing device). Both collective privacy preference models and app performance data can be used to make recommendations to users, as described further below. In an IoT context, app performance data may be replaced with data about the performance of different IoT resources under different possible configurations of user-configurable privacy settings.

[0047] FIG. **1A** is a block diagram of the computing device **114** according to various embodiments of the present invention. The computing device **114** comprises a processor (e.g., a CPU) and memory (not shown). The memory stores the code of one or more apps (software applications) **212** that the processor can execute to provide extended functionality to the computing device. Some, and indeed probably most, of the apps **212** request permissions to on-device functionality or data **140**, **141** and/or external functionality or data **150**, **151**. The on-device functionality **140** could be sub-systems of the computing device **114**, such as, a camera, a microphone, or a location determination sub-system (e.g., a GPS device, a location tracking sub-system combining data from multiple sensors such as GPS, Wi-Fi, accelerometer, etc.). The on-device data **141** could be information stored or otherwise determined by the computing device **114**. The information could be, for example, the user's contact information, the user's calendar information, the computing device's location (as determined by a location determination sub-system), health information, a device unique ID number, and many other types of information. Examples of permissions associated with external functionality and/or data **150**, **151** could include access to Facebook

functionality or data, or to external payment functionality or data, or to external smart home functionality or data, or to health and fitness data whether stored on a wearable device or in the cloud, as described above.

[0048] The computing device includes an operating system and, particularly germane to some embodiments of this invention, the operating system may enforce all or some of these permission settings. When a particular app **212** requests a particular functionality or data **140**, **141**, **150**, **151**, the operating system **210** may determine, based on the device's permissions settings, whether the app should be granted access to the functionality/data or not. If access should be granted, the operating system **210** can provide access to the functionality/data. In other embodiments, decisions about granting access to sensitive functionality or data may be the responsibility of Policy Enforcement Points (PEPs) capable of enforcing user-specific permission settings configured by users and/or their Privacy Assistants via APIs that the Privacy Assistants can discover in IoT Resource Registries, with the IoT registries advertising descriptions of IoT Resources including their user-configurable (or user-specific) privacy settings, as illustrated in FIG. **1B**. More details about PEPs for IoT uses can be found in A. Das, et al., "Assisting Users in a World Full of Cameras; A Privacy-aware Infrastructure for Computer Vision Applications," **2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)**, July 2017, which is incorporated herein by reference in its entirety.

[0049] As shown in FIG. **1A**, the computing device **114** may run the PPA **230**, which can be another app stored in the memory and executed by the processor. The PPA **230** as explained herein helps the user of the computing device **114** manage and/or otherwise configure privacy settings. This may include app permissions, privacy settings of the computing device **114** and user-configurable (user-specific) privacy settings of IoT resources **115**.

[0050] Also, the computing device **114** runs at least one user interface through which the user can interact with the computing device **114**. For example, the interface could comprise a display screen that shows text and graphics. The user could enter text via a keyboard, for example. The display screen could also comprise a touchscreen that recognizes touches of the display screen by the user. Alternatively or additionally, the computing device could include a speech-based interface where a speaker outputs audio and a microphone picks up the user's audible responses. This user interface functionality could be further enhanced with additional features including a haptic interface that taps the user's wrist to notify him or her about the presence of IoT resources with particular data collection or use practices, visor interface functionality, and other technologies that can be used to communicate with the user.

[0051] As shown in FIG. **1B**, and as described further herein, the Personalized Privacy Assistant **230** in other embodiments can also help users configure user-specific privacy settings associated with different external IoT resources such as IoT apps, IoT devices or IoT services. Examples of IoT apps include apps that allow employees in a given building to share their locations with one another, or an app to enable someone to remotely monitor their pet at home. Examples of IoT devices include a video camera, an HVAC system, and a smart speaker (e.g., Amazon Echo™). Examples of IoT services include a building's location tracking functionality whose output can be consumed by other IoT services and/or IoT apps such as the location sharing app mentioned above; a video monitoring system whose output can be consumed by other apps such as a video analytics system with algorithms to detect petty crimes; and an analytics system to monitor the particular products a consumer in a supermarket is looking at. In this particular configuration, the IoT resources may be discovered by the Personalized Privacy Assistant **230** using discovery protocols that rely on IoT resources being advertised in IoT resource registries **200**, which can themselves be discovered by the Personalized Privacy Assistant **230** based on the location of the user (e.g. discovering an IoT Resource Registry associated with a building as the user of the Personalized Privacy Assistant approaches that building). The IoT registries **200** may be queried based on the location of the user, with entries in the registries corresponding to IoT resources (e.g., IoT apps, IoT

devices or IoT services) that are advertised as being deployed in different areas. Registry entries for IoT resources include other relevant attributes, such as attributes describing the IoT resource's data collection and use practices (e.g., what data is being collected, how long it is retained, whether it is aggregated or anonymized, for what purpose it is collected, which third parties it might be shared with, if any, etc.) and the user-specific settings exposed to users, if any (e.g., discovering in an IoT Registry an entry corresponding to a camera monitoring system in a building with a description of user-specific opt-in setting that authorizes the system to apply facial recognition and scene recognition functionality to the video streams it captures). The user-specific settings are advertised along with APIs that can be used by Personalized Privacy Assistants to communicate directly with these IoT resources or with policy enforcement functionality associated with these IoT resources to configure these user-specific settings (e.g., API for a user to opt into facial recognition for the camera monitoring system via his or her Personalized Privacy Assistant).

[0052] In various embodiments, the IoT Resource Registries **200** can be discovered by the Personal Privacy Assistants **230** based on the location of the user (e.g. using the location of his/her computing device **114** as a proxy for his or her location). This can be done in several possible ways. In some embodiments, different areas may be equipped with short range beacons advertising the registries (e.g., advertisements that include URIs where the registries can be accessed/queried). The user's computing device **114** can receive these advertisements wirelessly from the beacons and the PPA **230** can use the URL in the advertisement to access/query the registry **200**. In some embodiments, registries may be associated with "areas of coverage" and may themselves be listed in one or more directories **123** that can be accessed by the Personal Privacy Assistants **230** (e.g., the Personal Privacy Assistants **230** may have an initial list of predefined directories **123** it knows about, or it may discover these directories **123** over time or have its list of directories **123** updated over time). In some embodiments, when querying a directory **123**, the Privacy Assistant **230** can provide the user's location (e.g. location of the computing device **114** on which the Privacy Assistant **230** is running, if it is running on a mobile computing device carried by the user). Based on that location, the directory **123** can return a list of registries **200** relevant to the user's location based on the areas of coverage of the registries listed in the directory **123**. The Privacy Assistant **230** can then query each of these registries **200**, optionally providing the user's location again as part of the queries, and receive in return a list of IoT resources listed in these registries **200** with the list of IoT resources optionally being limited to IoT resources whose areas of coverage are in the vicinity of the user's current location (e.g., within some distance, or satisfying some other criteria indicative of whether the user is likely to have his or her data collected by these devices such as being within the range of a system of cameras). That way, the PPA **230** can control the collection and use by the IoT resources of sensitive user data and/or control access to sensitive functionality of the IoT resource (e.g., functionality provided by a camera, location sensor, microphone, biometric sensor, payment or point-of-sale system, etc.).

[0053] FIG. **3** illustrates the various functions of the PPA **230** according to various embodiments. As described herein, the PPA **230** can store a local copy of the collective privacy preference model developed by the data center (see FIG. **2**), and the PPA **230** can be responsible for instantiating the collective model for a particular user (the computing device's user) based on data collected about that user (e.g. apps installed on the user's phone, user answers to personalized questions, etc.). The PPA can also be responsible for further refining that user's individual privacy preference model based on additional information collected from the user such as the particular permission recommendations the user accepts, rejects or modifies as well as additional permissions the user configures over time and possibly answers to additional personalized questions asked over time to the user.

[0054] In some embodiments, the PPA can report some or all of the user collected information to the data center **106** whether (a) to receive a privacy profile corresponding to the cluster identified for the user and optionally to the specific apps installed on the user's device, (b) to just allow the

data center to update its models, or to support some other activities, as determined by the particular way in which functionality is distributed between the PPA and the data center—different configurations offer different advantages in terms of local processing requirements, communication lags, responsiveness, and more. In other embodiments, as shown in FIG. 3, the PPA can store a local copy of the collective privacy preference model in order to generate an individual privacy preference model based on the user data. The PPA as described herein can also generate questions for the user to elicit their privacy preferences, provide privacy nudges to the user to sharpen models of their preferences, motivate them to review their settings, help determine the cluster that best matches the user's preferences, and determine, present, confirm and implement permission recommendations.

[0055] In various embodiments, the PPA **230** uses information about the apps **212** installed on a user's computing device **114** to elicit the user's privacy preferences, such as (in one embodiment) whether the user feels comfortable or not granting a particular permission to apps in a given category (e.g. game apps, utility apps, social networking apps). This may come in the form of questions customized to reflect the particular apps in this category that the user has installed on his or her device. These questions can be further qualified according to the particular purpose(s) for which apps in a given category may request a given permission (e.g., determining whether the user is generally comfortable granting access to his fine location gaming apps category for the purpose of the apps being able to operate, for the apps to share the resulting information with advertising networks that may in turn use this information to better target the user, or for the apps to share this information with analytics companies that collect extensive profiles about users, or determining whether the user is comfortable granting shopping apps access to camera functionality for the purpose of scanning bar codes, etc.). Based on the information collected from the user, the PPA can offer the user personalized recommendations on how to configure associated permission settings. The PPA can use interactive profile assignment dialogs for this purpose. The dialogs can be dynamically generated by the PPA based on, for example, decision trees generated to identify a cluster that best matches a given user. In one embodiment, the dialogs consist of questions asked to the user by the PPA, as it traverses the decision tree, with questions at each node in the tree taking into account information such as the particular apps the user has installed on his or her computing device and the answers he or she has provided to earlier questions. The dialogs may be used to match users to clusters and corresponding privacy profiles that best align with their preferences. The profile identified for a given user can be the basis for an individual privacy preference model for that user (in embodiments of the PPA that rely on privacy profiles), which the PPA can further refine over time based on additional interactions with the user. The determined profile can be used to provide the user recommendations on which app permissions to grant or deny and more generally how to configure different user-specific privacy settings. In other embodiments the dialogs with users may be based on other collective privacy preference models and may rely on other machine learning techniques such as content filtering techniques or some of the other statistical and machine learning techniques discussed herein. In some embodiments, the PPA gives the user the option to accept multiple recommended settings at once and possibly also the ability to modify one or more recommendations, as needed.

[0056] FIG. 12 is a flow chart of a sample process by which the preference models can be built and used in the PPA. Process **2100** involves generating the models and includes collecting permission settings from the test subjects at step **2120** as well as possibly other relevant data (e.g. log of permission setting changes for the user) and forwarding these test subject privacy data to a test subject privacy database, typically stored on a server. In some embodiments, prior to collection of the data or in parallel with it, some or all of the test subjects can be exposed to privacy nudges at step **2110**. “Privacy nudges” can be considered as communications with a test subject (e.g., through pop up messages or other ways of catching the test subject's attention) aimed at motivating the test subject to reflect on certain privacy decisions (or privacy choices) they have made or need to make,

whether implicitly or explicitly (e.g. decisions associated with permission settings or other user-specific privacy settings). A privacy nudge can highlight potential risks to a user that the user may not be aware of or may have underestimated, such as the number of times the test subject shared his/her location, how that information is used, who it is shared with, what might be inferred from it and more. Privacy nudges are designed to overcome cognitive and behavioral biases such as user's tendency to heavily discount potential risks associated with the collection of their data and the different ways in which this data could potentially be used, and behavioral biases such as user's tendency to not engage with default privacy settings as doing so is often perceived as a distraction from the primary task in which the user is engaged (e.g. posting an update on a social network, responding to an email, completing a purchase online). More details about privacy nudges can be found in H. Almuhimedi et al., "Your Location has been Shared 5,398 Times! A Field Study on Mobile App Privacy Nudging," 33rd Annual ACM Conference on Human Factors in Computing Systems (CHI), pp. 787-796 (2015), which is incorporated herein by reference.

[0057] FIG. 4 shows an embodiment of a privacy nudge **400** graphically rendered on, for example, the user's computing device **114** according to various embodiments of the present invention. In the illustrated example, the privacy nudge indicates the total number of instances that an app was granted access to data or functionality controlled by a permission, namely in the illustrated example, the user's location. Moreover, an app component **404** can display the number of instances that particular apps used a permitted privacy-related permission, such as accessing location data, as in the illustrated example. In addition, in this particular instance, a purpose component **406** of the privacy nudge **400** indicates the probable purpose for which the accessed location data is used, such as targeted advertising or consumer tracking and profiling. Such purpose information might further motivate a user to review his or her settings, as it might come as a surprise or might reveal to the user data practices that he or she might not feel comfortable with. In this particular instance, purpose associated with each permission request was derived from an analysis of the app's code, looking in particular at whether requests for particular permissions originate from third party libraries or not, and inferring the likely purpose(s) for accessing the permission accordingly (e.g., permission accessed to support the app's core functionality when first party code is responsible for the request, permission accessed to share information with a social networking site when the permission request originates from a third party library associated with a social networking site, with an analytics company when it originates from a third party library associated with an analytics company, with an advertising network when it originates from a library associated with an advertising network, etc.), as detailed in J. Lin et al., "Modeling users mobile app privacy preferences: Restoring usability in a sea of permission settings," 2014 Symposium On Usable Privacy and Security (2015), which is incorporated herein by reference. In other embodiments such information can be inferred from descriptions of the apps, from information directly provided by the app's developer, from analysis of the app's behavior and the traffic it generates, or even from user reviews.

[0058] After collection of the test subjects' permission settings at step **2120**, at least one, but possibly multiple, collective privacy preference models may be derived from the test subjects' permission setting data. The model(s) may be derived using machine learning or statistical analysis techniques. Examples of relevant machine learning and statistical analysis techniques that can be used to develop collective privacy preference models include: clustering techniques, collaborative filtering techniques, content-based filtering techniques, logistic regression techniques, support vector machine techniques, Bayesian inference techniques, decision tree learning techniques, and more, including ensemble methods that combine multiple models. Preference models can also be learned or improved using deep learning techniques, including techniques relying on the use of ad hoc or general purpose foundation models. Other machine learning techniques, including other unsupervised learning techniques can also contribute to further enriching collective privacy preference models and individual privacy preference models. Privacy preference models might also

be produced through the manual refinement of rules or through dialogues with users.

[0059] In one embodiment, collective privacy preference models can come in the form of a number of privacy preference profiles obtained by clustering users with similar privacy preferences and identifying profiles of permission settings that users in a given cluster strongly agree on. The PPA can use these profiles to identify for each user the cluster (and profile) that best matches his/her privacy preferences. The privacy preference profiles can comprise, in one embodiment, collections of permission settings that test subjects in a given cluster strongly agree on (e.g. a threshold percentage of test subjects in the cluster concur on granting a given permission to a given category of apps for a given purpose, or concur on denying a given permission to a given category of apps, or more generally settings identified for different clusters of users using machine learning and statistical techniques such as Support Vector Machines, Random Forests, or other available machine learning and statistical techniques). A clustering technique for deriving the profiles is described in more detail herein. Different possible features can be used to identify clusters of users and associated privacy profiles, as well as build other types of privacy preference models. This includes building finer models that distinguish between individual apps (especially in the case of particularly popular apps, for which ample data from test subjects is likely to be available), building models that distinguish between finer types of purposes, building models that distinguish between different categories of third parties with which data might be shared or even individual third parties, building models that take into account additional information such as the retention policy associated with the collection of data, whether the data is aggregated or anonymized and more.

[0060] Process **2200** shows the general process in which the model is used to configure permission settings for a user, i.e., a user of computing device **114**. At step **2210**, information about the user is collected. Various techniques and sources for collecting the user's information are described herein, and can involve, in one embodiment, answers to fixed questions, or to dynamically generated questions personalized for the user, e.g. based on the specific apps the user has downloaded on his or her computing device. These personalized questions can be presented to the user via an interface of the computing device (e.g., a text/display interface and/or a speech-based interface).

[0061] FIGS. **7A-C** show examples. In these examples, the PPA **230** presents a display to the user asking the user how he/she feels about apps in a certain app category accessing a particular functionality or data **140**, **141**, **150**, **151**. In FIG. **7A** the app category is social apps and the requested sensitive data is the user's location; in FIG. **7B** the app category is also social apps but now the requested functionality is the user's computing device's camera; and in FIG. **7C** the app category is finance apps and the requested data is location. At the bottom of each display, the user can input their response. The user's responses can be used to make the permission recommendations to the user as described herein.

[0062] In various embodiments, this information (e.g., the user's response to the question(s) or other applicable user information) is transmitted to the data center (such as from the user's computing device or other sources of such user information) so that, at step **2220**, recommended permission settings for the user can be identified by the servers of the data center based on the user information collected at step **2210**. In other embodiments, the parts of the collective privacy preference models required to assign a cluster to a user can be stored on the computing device itself (**114**) and the computation can be performed locally on the computing device, at which point the corresponding privacy profiles can be retrieved from the data center (e.g. by identifying to the data server the particular cluster determined to be relevant and the list of apps installed on the user's computing device). In yet other embodiments, the privacy profiles for each cluster can also be stored on the user computing device, in which case the entire process can be performed locally. In one embodiment, the user can be assigned a particular cluster (the cluster that most closely matches information collected about the user) and the recommended permission settings can be based on the permission settings found in the privacy profile associated with the cluster assigned to the user,

such as illustrated in the examples of FIGS. **14-16** described below. In some embodiments privacy profiles may include recommendations at the level of entire categories of apps, while in other embodiments some or all recommended settings may be at the level of individual apps (e.g., particularly popular apps for which sufficient test subjects data is available and/or for which users might feel differently about granting or denying some permissions, possibly because they trust these apps more or for some other reason). In yet other embodiments, as already described herein, recommendations can be filtered based on operational data (or app performance data) about the apps, such data indicating whether denying particular permissions to particular apps seem to be impacting performance (of the app itself, of the device or possibly of other apps). In some embodiments this operational data may also be stored on the computing device **114**. In other embodiments it might be stored in the data center **106** with only data relevant to those apps installed on the computing device being retrieved by the computing device from the data center to determine which specific permission setting to recommend to the user.

[0063] Specifically, the PPA can look at each app installed on the user's device, identify the category corresponding to this app (in the case of a privacy preference profile where recommended settings are organized by app categories—other models are possible), and identify in the profile recommended settings for the permissions requested by this app, if there are any. At step **2230**, the recommended permission settings are presented to the user of the computing device and at step **2240** the permission settings can be configured based on feedback from the user, such as the user accepting, rejecting or modifying the recommended permission settings.

[0064] FIG. **8A-C** show illustrative displays according to embodiments of the present invention of how the user could provide feedback. First, as shown in FIG. **8A**, the user could be presented with a screen where permissions are organized by types of permissions (e.g., location, contacts, messages, etc.) with an identification of the number of apps that request each type of permission. The user could then select one type of permission, which then causes the display to list the apps that request that type of permission, along with any recommendation by the PPA to possibly deny some of these permissions, if they are not yet denied. The reader will appreciate that a number of ways of presenting recommended permission settings are possible. In this particular embodiment, the user can then go through the recommended setting changes (e.g., recommendation to change the location permission requested by the PayPal app from “grant” to “deny” in FIG. **8B**, or to change the (access) contacts permission from “grant” to “deny” for the CitiMobile app in FIG. **8C**). In this particular embodiment, the user can accept or ignore each recommended setting change. In doing so, the user provides feedback to the PPA. This feedback can in turn be used to refine the user's Individual Privacy Preference model (step **2310** in FIG. **12**). In this particular embodiment, the user can review each permission for each individual app as shown in the examples of FIGS. **8B** (app requesting location) and **8C** (apps requesting contacts), decide whether to accept or ignore different recommended settings, modify settings for which the PPA does not provide any recommendations, or modify earlier setting changes, including changes that may have been prompted by the PPA's recommendations.

[0065] In some embodiments, the recommendations display provided by the interface to the user may also provide the user with the option of requesting an explanation of the recommendation. One such embodiment is shown in FIG. **5**. In the illustrated embodiment, several recommendations are made by the system. Each recommendation has an input means, such as a slide switch, that allows the user to allow or deny a given permission to an app. There is also, in the illustrated exemplary embodiment, a question mark (“?”) for each recommendation. Of course, in other embodiments, other icons or symbols could be used to provide a means for the user to request an explanation for the corresponding recommendation. In the illustrated embodiment, the user can click the question mark, and in response the PPA can display an explanation of why a particular setting is recommended. In the example of FIG. **5**, the user clicked the question mark for the Snapchat location recommendation and the PPA displayed why the PPA recommended denying that

permission. That way, the user can make more informed decisions with regard to the recommendations. Such explanations can be generated in several possible ways. One approach is to identify the conditions that triggered a given recommendation. For instance, in the embodiment displayed in FIG. 5, the user has been matched with a cluster in which the majority of users do not feel comfortable disclosing their location to social apps when their location information is potentially used for “consumer tracking and profiling”. Accordingly, the explanation shown to the user identifies these conditions as the conditions that triggered the corresponding recommendation. Other possible approaches for generating recommendations include taking into account information gain considerations that can be used to identify those factors most likely to have contributed to a given recommendation. The number of such factors can be limited by a threshold to ensure that the recommendations are easy to understand. The threshold could possibly be adjusted (e.g. increased) in response to additional requests from the user or some other type of interaction with the user indicating that he or she would like additional details. In other embodiments, the PPA might generate explanations that refer back directly to information collected from the user such as answers provided earlier by the user to some questions (e.g. “You told us you didn't feel comfortable sharing you location with Facebook and TikTok. In general, people who don't feel comfortable sharing their location with Facebook and TikTok, also do not feel comfortable sharing it with Snapchat”). In yet some other embodiments, explanations could also include an opportunity for the user to indicate that he or she feels differently about a particular recommendation (e.g., “If you don't feel the same way about Snapchat, click here”). This could in turn be used to trigger a dialog intended to further clarify the user's privacy preferences and better discriminate between situations when he or she feels comfortable granting a permission and those when he or she feels uncomfortable doing so.

[0066] In some embodiments, recommended settings can be presented in bulk with the user having the ability to review them and decide individually which recommendation to accept, reject, or modify, as shown in the example of FIGS. 8A-C. In other embodiments, the recommendations can be presented more incrementally to the user, such as when a user launches a particular app on his computing device, or when the user has just completed the installation of a new app, as shown in the example of FIG. 9. In yet other embodiments, recommendations might be bundled, with the user only having the option to accept or reject multiple recommendations at a time. In yet other embodiments, users may also have the option of modifying recommendations (e.g., when permissions are non-binary such as a permission where a user can choose to modulate the level of granularity at which a permission is granted such as access to location at different levels of granularity).

[0067] In yet other embodiments, additional personalized questions can be generated when the user accepts, rejects or modifies a recommended setting. FIG. 10 illustrates such a situation, namely User 1 denies Permission 3 to App 21, despite the fact that User 1 belongs to cluster 1 and that the privacy profile for cluster 1 recommends granting permission 3 to apps in App Category2. At this point, in some embodiments, the PPA can generate a personalized question to see whether this rejection of a recommendation can be used to infer a more general privacy preference for User 1, namely by asking something like, “In general, do you feel uncomfortable granting Permission 3 to apps in App Category 2”, or, as is assumed in FIG. 10, “In general, do you feel uncomfortable granting Permission 3 to any app?” In the particular instance illustrated in FIG. 10, User 1 answers “yes” to the latter question. This in turn results in the system updating the individual privacy preference model for User 1 and noting that, in contrast to the general profile for users in Cluster 1 (part of the collective privacy preference model), User 1 wants to systematically deny Permission 3 to all apps, as denoted by the two asterisks next to “Deny” for the entry for Permission 3 in the rows corresponding to App Category 2 and 3.

[0068] Dialogs of this nature can be initiated by the Privacy Assistant under a number of possible situations. This includes the time when a user installs another app or is about to first come into

contact with a new IoT Resource that has just been discovered by his privacy assistant. This also includes situations where the privacy assistant sends a nudge to the user (e.g. to motivate the user to revisit prior privacy decisions and see whether his or her preferences have not possibly changed). This can include situations where a user modifies a permission setting he or she had selected earlier. The question(s) in this case can focus on understanding the change and trying to elicit the scope associated with this change (e.g. is it specific to the context in which the user is, is it specific to the app or IoT resource with which the user is interacting, does it have to do with the category of app or IoT resource, the purpose for which the permission is requested, the retention period associated with the data being collected, the nature of the permission, or the potential uses of the information being collected). Different possible scopes can lead to different questions, with these questions potentially being scripted or being driven by data collected from users and privacy preference models that can help identify those questions that are likely to be most discriminative (e.g., maximizing information gain to minimize the number of questions that the user has to answer, for instance by learning decision trees associated with different situations such as changes of different types of permissions or different patterns of change in permissions associated with different situations).

[0069] Asking the user a follow up question, such as a question based on his feedback on recommended permission settings, as outlined above, corresponds to step **2310** in the process outlined in FIG. **12**. It can take place after permissions have been configured or possibly as part of a dialog with the user that can take place while permission settings are recommended to the user (namely as part of step **2230**). The present invention envisions any number of ways of interleaving these interactions with the user.

[0070] Once the permission settings are configured, whether automatically or after checking with the user, as outlined above, they start to be enforced of the computing device's operating system.

[0071] Process **2300** shows a general process for further refining a user's Individual Privacy Preference model. At step **2310**, additional information about the user can be collected. It can take the form of a question that directly builds on the feedback of the user on permission settings recommended to him/her, as illustrated above. It can also be triggered by the user downloading a new app on his computing device and manually configuring one or more permissions requested by this app. It could also be triggered by a number of other events such as the user deciding to manually review some of his or her permission settings, thereby opening the door to another dialog with the PPA and an opportunity to learn more about the user's privacy preferences. This information will typically be stored locally by the PPA and may be fed into the PPA's local machine learning functionality to derive additional preferences for the user and/or generate additional personalized questions for the user and further elicit information about his or her preferences. This information can also be forwarded to the data center for incorporation in the collective privacy preference models. Step **2320** represents a step where the additional information collected from the user is used to refine the user's individual privacy preference model. As the user's individual privacy preference model is refined, additional recommendations (step **2330**) may be identified and presented to the user, whether in bulk or in a more incremental manner (e.g. as the user launches or downloads apps). In some situations, they may simply be automatically implemented (e.g. if confidence in the recommendation is sufficiently high).

[0072] In some embodiments, permission recommendations are filtered based on app performance data, as illustrated in step **2225** in FIG. **13**. An example of app performance data is shown in FIG. **11**. App performance data will typically be collected by a data center/server **106** (see FIG. **2**) and can include data such as whether the app works when a given permission is denied or whether it crashes or loses some of its functionality. This information may be directly collected from the devices of test subjects, or it may be obtained by systematically testing each app while selectively denying each of its requested permissions. This information could also be obtained from or supplemented with static code analysis techniques that predict what happens when a permission is

not granted. In some embodiments, some of this data could be provided directly by app developers. Another valuable source of app performance data can be obtained by mining reviews of apps by users, focusing on comments that pertain to users who tried to deny some permissions. As users download new apps on their computing device, their PPA can fetch performance data that pertain to these apps and store it locally to help filter recommended permission settings it generates. As discussed earlier, different embodiments might store data in different places (e.g. whether on the user device or in a data center).

[0073] User privacy preference models may be limited to the granting or denying of permissions to different apps. FIG. 14 illustrates an example for two different clusters—Cluster 1 and Cluster 2. The two tables in FIG. 14 show recommendations for people in each cluster, namely preferences common to a significant majority of test subjects assigned to that cluster. Permissions 1, 2 and 3 could be different types of permissions that apps frequently request. In the illustrated example, for apps in a first category (App Category 1), users in Cluster 1 have a sufficiently strong tendency (e.g., above a threshold confidence level, or as determined via some machine learning technique) to grant Permission 1 and deny Permission 3, but there is no strong agreement (or recommendation) for Permission 2. On the other hand, users in Cluster 2 collectively show sufficiently strong agreement on denying Permissions 1, 2 and 3. Thus, if a new user was determined to match most closely with users in Cluster 2, the permissions recommendations for the new user would be to deny Permissions 1, 2 and 3 for apps in App Category 1, and so on as illustrated in FIG. 14. In some models, users could be assigned to multiple clusters. In other models (e.g. collaborative filtering models or other models discussed elsewhere in this invention), recommendations on whether to grant or deny different permissions to different apps or different app categories under different conditions (e.g. different purposes) may be computed directly without having to assign users to clusters. Recommendations may include recommendations to deny or grant a permission to a particular app, or a category of apps; they may include recommendations to prompt the user to configure a setting; they may include differentiating between purposes associated with a permission or privacy preference setting, or other relevant data collection and/or use attribute associated with a given user specific privacy setting. They may also include recommendations to modify the granularity or resolution at which some data or functionality can be accessed. More generally, privacy preference models are not limited to recommending privacy settings and but may also include the learning of models of when to notify users about different types of data collection and use practices, as discussed elsewhere in the present invention.

[0074] In the example of FIG. 14, the recommendations were based on the app category and the permission requested. In other embodiment, the models might differentiate between the different possible purposes for which an app might request a given permission, as described in connection with FIG. 15. For example, a user might be more likely to grant a permission for an app where the requested functionality/data is needed for the core operation of the app, and less likely to grant the permission where the requested functionality/data is not needed for the core operation (e.g., advertising). FIG. 15 is a table showing how, for one cluster of users, the recommendations can change based on the purpose of the requested permission. In the illustrated example, it is assumed that there are three possible purposes for each of two permissions (not every app in a given category has to necessarily request each permission or request a permission for each possible purpose—this will typically vary from one app to another). As shown in this example, people in the cluster tend to grant Permission 1 for Purpose 1 for apps in App Category 1, but there is no consensus (and no recommendation) for Purposes 2 and 3 of Permission 1 for App Category 1, and so on as shown in the example of FIG. 15. Thus, the permission recommendations to a new user that was classified to this cluster would be to grant Permission 1 for App Category 1 for Purpose 1, and deny Permission 1 for App Categories 2 and 3 for Purpose 1, and so on as shown in the example of FIG. 15.

[0075] The models may further distinguish between different contexts or the values of different

contextual attributes, with some users feeling comfortable granting some permissions under some contexts but not under others. For instance, a user might feel comfortable granting access to his or her location when he or she is at the office or in a public place, but not when he or she is at home or at a bar. In this example location is a contextual attribute whose value determines when the user feels comfortable granting a permission. Research shows that many users have such context-sensitive privacy preferences. In the absence of a PPA and of permission settings that capture contextual attributes such as these, a user would have to manually toggle corresponding permission settings as his or her context changes. In some embodiments, a PPA can take care of automatically modifying the user's settings, based on his or her preferences and the value of relevant contextual attributes. An illustrative example of such context-sensitive permission settings for a cluster is shown in FIG. 16. The example shown in FIG. 16 includes some permission recommendations that were not used in the examples of FIGS. 14 and 15. For example, one illustrated recommendation is "grant low fidelity access" (e.g., for Permission 1, Context 3 for App Category 1). When a permission request for functionality/data has this setting, the full capabilities of a sub-system of the computing device or the highest resolution information of the computing device is not provided. Instead, some diminished capability is provided, such as location only to within a certain distance (e.g., 1 mile) or lower resolution camera images or sound captured by the computing device's microphone, etc. Another different setting is "Falsify" (Permission 1, Context 2 for App Category 3). When a permission request has this setting, the app is granted false or spoofed information, such as false location information, false contact information, etc. Other permissions are "grant restricted access," such that, for example, the app's possible use of the functionality/data granted by the permission is restricted.

[0076] A field study conducted by the inventors is described in the incorporated provisional application.

[0077] For the purpose of accurately capturing users' privacy preferences from their privacy settings, in one embodiment the inventors assume that users (or test subjects) are comfortable with a restrictive permission setting they chose, if they have subsequently been observed keeping the restrictive setting (i.e. not changing it back to a permissive setting) over a sufficiently long period of time. To increase users' (or test subjects') awareness and engagement, and to motivate them to review their permission settings, in one embodiment, the inventors developed a PPA by making a number of modifications and enhancements to the Android permission manager App Ops, as described below. It should be clear that similar functionality could be added to any number of permission managers and/or included in other forms of PPA (e.g. a PPA to help users configure user-specific privacy settings associated with IoT Resources).

[0078] As mentioned herein, in one embodiment privacy preference modeling may use hierarchical clustering (or other clustering techniques) to group like-minded users and generate associated privacy profiles. In various embodiments, the clustering can use features that include (1) the likely purpose of the permission request as determined based on available datasets that collect such information (e.g., as described by the app developer, as identified using static and/or dynamic analysis of the app's code, as obtained by mining user reviews of the app, or using any other relevant source of information available to infer the purpose associated with a given permission request) and (2) the category of the requesting app from the Google Play store or any other suitable app store (e.g., game, social, finance, news, etc.) or any other meaningful way of organizing apps into categories, (3) the context in which a particular permission is being requested (e.g. whether permissions are restricted to times when the app is actively used by the user or whether they also include background processing, other relevant contextual attributes that might restrict conditions when a permission is being used or requested such as where the user is, what the user is doing, time of day, etc.), (4) other relevant attributes characterizing the collection and/or use of the data or functionality controlled by a given permission (e.g. retention period associated with the data being collected, level of granularity or aggregation at which functionality or data is being accessed, who

the data might be shared with, etc.). App categories, rather than individual apps, can be used as features to reduce overfitting caused by less popular apps and limited training samples. Some models may also combine app categories with individual apps (e.g. popular apps may yield sufficient data from test subjects to warrant being treated separately, and/or may elicit different privacy preferences from users compared to similar yet less popular apps in the same category). User privacy profiles can be built based on training data collected from test subjects (or more generally from users), by applying, in one embodiment, clustering techniques on the feature vectors capturing data collected from test subjects or users (e.g. which permissions they are willing to grant to different apps or categories of apps and other relevant features such as those identified above). More details about hierarchical clustering and some of the other clustering and machine learning techniques that can be used to derive privacy preference profiles as applied according to some embodiments of the present invention can be found in the scikit-learn User Guide, release 0.19.1, Nov. 21, 2017, available at scikit-learn.org/stable/_downloads/scikit-learn-docs.pdf, as well as at scikit-learn.org/stable/user_guide.html, which are incorporated herein by reference. Other relevant clustering techniques include deep learning based clustering techniques and LLM-based clustering techniques, such as described in V. Viswanathan et al., "Large Language Models Enable Few-Shot Clustering," Transactions of the Association for Computational Linguistics (2024) 12:321-333 ([direct.mit.edu/tac/article/doi/10.1162/tac_a_00648/120476/Large-Language-Models-Enable-Few-Shot-Clustering](https://direct.mit.edu/tac/article/doi/10.1162/tac/article/doi/10.1162/tac_a_00648/120476/Large-Language-Models-Enable-Few-Shot-Clustering)), which is incorporated herein by reference. In other embodiments, other machine learning and statistical analysis methods can be used to create the model and derive privacy profiles (e.g. deep learning techniques), as discussed elsewhere in the present invention.

[0079] In one embodiment, the inventors quantified each user's preferences as a three-dimensional tensor of aggregated preferences of (app category, permission, purpose). For each cell, the value was defined as the tendency of the user to allow or deny permissions requested by apps from a specific category for a corresponding purpose: from -1 (minus one) (100% deny) to 1 (plus one) (100% allow), and N/A if the user's settings data was not available for a cell. To estimate similarities among participants' feature tensors, missing values were imputed in the tensors. In one embodiment, in order to impute without biasing any dimension, weighted PARAFAC Tensor factorization was applied. A weight of 1 was put on all known data cells and 0 weight on unknown data cells in the tensor. Thus, the overall error of the imputed tensor in Frobenius norm was optimized using only the values known from the data. Using the users' feature vectors reshaped from the imputed tensor, the user profiles were built by applying hierarchical clustering on the feature vectors. Hierarchical clustering provides an advantage that it is not sensitive to the size or density of clusters and allows non-Euclidean distances. The inventors envision many other variations of this embodiment, including variations that include tensor entries for individual apps rather than entire app categories or even mixed tensors with some entries corresponding to collections of apps in a given category and some entries corresponding to individual apps (e.g., particularly popular apps for which data is available from a large enough collection of users). Some embodiments may fill unknown data cells according to a number of other possible schemes. Some embodiments may rely on other distance metrics and other connectivity-based clustering techniques. Other embodiments may rely on other clustering techniques (e.g., other connectivity-based clustering techniques, centroid-based clustering techniques such as k-means clustering techniques, distribution-based clustering techniques such as Gaussian mixture model techniques, density-based clustering techniques such as DBSCAN or OPTICS, canopy pre-clustering techniques, neural network techniques and more) or other combinations of clustering techniques. And, as already discussed earlier, recommendations may also be derived using other machine learning and statistical analysis techniques that do not rely on clustering.

[0080] In one possible embodiment, in order to assign new users to the generated privacy profiles, the PPA can ask the users a small number of questions about their privacy preferences. Some or all of these questions can be tailored to the particular apps installed by the user on his or her

computing device and relevant features associated with these apps (e.g., as captured in the collective privacy preference model). They can also be tailored based on other information available about the user such as information identified as relevant to identifying clusters that best match users (e.g. features in the privacy model), information that can help personalize questions to reflect relevant attributes of the user and can help make these questions more relevant to the user (e.g. demographic information, home address, where the user works, information about the user's friends, etc.), as well as information related to the behavior of the apps installed by the user on his or her computing device (e.g. how often different apps have requested access to some permissions, for what purpose, etc.). This information can also be used to eliminate questions that are not relevant to a user. For instance, a generic decision tree to assign users to clusters might include a question about whether users are comfortable sharing their contacts lists with social apps for advertising purpose. If a given user does not have any social apps on his or her smartphone, the PPA could decide to build a decision tree for that user that does not include questions about social apps. In one embodiment, the inventors have used a C4.5 decision tree (more details about decision trees used according to embodiments of the present invention can be found in J. R. Quinlan. *C4.5: programs for machine learning*. Elsevier, 2014, which is incorporated herein by reference) on the set of questions applicable to a particular user (e.g. based on the apps a given user actually has on his or her computing device, the permissions requested by these apps and the purposes for which these permissions are being requested), and generate an ordered list of questions. Other decision tree algorithms could readily be used to generate these questions (e.g. ID3, C5, etc.), as well as other statistical classification algorithms. In one embodiment, users are asked a maximum of five questions to be assigned to a profile so as to limit the time investment of the user. The set of questions can be dynamically personalized for each user, so that the questions can be contextualized using the apps each user has installed on their computing devices (e.g., smartphones). They can also be dynamically contextualized to reflect other relevant attributes of a user's context such as where the user is, whether the app is actively being used versus running in the background, etc. (e.g., "In general, when at home, do you feel comfortable disclosing your location to utility apps," or "When you declined access to location to the Groupon app, was it because it was running in the background or because you generally don't want to share your location with Groupon?").

[0081] Collective Privacy Preference Profiles for different Clusters of Users (e.g., FIG. 10) can be derived by looking at the settings collected from users (or test subjects) for different permissions and different app categories (e.g. as illustrated in FIG. 6). More generally, this data may also be organized to differentiate between different types of purposes or any number of other relevant attributes as already discussed earlier in this invention (e.g. contextual attributes, data retention attributes, etc.) when such attributes are identified as influencing people's preferences or when permission settings can be configured differently for different values of these attributes. Collective privacy preference profiles can also be built for other user-configurable privacy settings as well as for privacy notification preferences, as also discussed earlier already. Recommended settings associated with these privacy preference profiles can be derived in several possible ways. In some embodiments, recommended settings for users in a given cluster will be identified when a sufficiently large fraction of users (or test subjects) in a given cluster are determined to have a preference for the same setting (e.g. a fraction of users above a given threshold such as 75% of users in the cluster have indicated they want to deny apps in a given category access to a given permission for a given purpose). In such embodiments, settings for which there is insufficient agreement among users (or test subjects) in the given cluster may simply not have any recommendation associated with them. In other embodiments, multiple thresholds might be used, including a threshold above which a recommendation will be automatically enacted (e.g. say a threshold of 90% and a sufficiently large number of data points from users or test subjects), and one where the recommendation will be used to prompt the user and see whether he or she agrees

with the recommended setting (e.g. say a threshold of 70%). In other embodiments, recommendations may be derived using a number of available machine learning techniques. For instance, in one embodiment a support vector machine (SVM) classifier (more details about techniques relevant to different possible embodiments of the present invention can be found in R.-E. Fan, K.-W. Chang, C.-J. Hsieh, X.-R. Wang, and C.-J. Lin, "Liblinear: A library for large linear classification," *The Journal of Machine Learning Research*, 9:1871-1874, 2008, which is incorporated herein by reference) can be trained using the permission settings (or privacy-configurable settings or privacy notification settings) that were collected from users (or test subjects) in a given cluster. Other embodiments may rely on non-linear machine learning models as well. Based on the resulting privacy profiles and, more generally, based on these recommendation models, the PPA app may use the specific features associated with a given user to generate recommendations for privacy settings for that user. The features that are included can be, for example, the user's assigned profile, a particular app category, a particular permission, and one or more purposes, or any other relevant collection of such features. For instance, in one embodiment, such a collection of features might be provided as a collection of tuples with each tuple including an app category, a permission, and a purpose. In some embodiments, recommendations may also take into account relevant elements of the user's context (e.g., granting access to the user's location when the user is in a public space, and denying such access when the user is not). In some embodiments the recommendations may come in the form of an "allow", "deny", or "no recommendation" answer for each of these tuples. More complex recommendations may include recommendations to provide limited or even obfuscated access to a given permission, as discussed elsewhere in the present invention, and more generally recommendations associated with non-binary privacy permission controls (e.g., multiple levels of granularity at which access to sensitive data might be granted). In some embodiments, if the set of permission settings available is more limited than what the model is able to predict, such as in the more limited permission model supported in Android 6, which does not allow to differentiate between different types of purposes for which a permission might be requested, the recommendations might come in the form of coarser recommendations, namely recommendations that take into account the different preferences users in a given cluster might have for granting a permission for combinations of possible purposes. In yet other embodiments, recommendations may also take into account the impact of denying a permission on the performance of an app, or the performance of the computing device, as discussed elsewhere in this invention.

[0082] In one embodiment where seven different clusters of users have been identified, FIG. 6 shows aggregated permission preferences of users in each cluster organized by app categories. This figure illustrates the typical diversity in privacy preferences associated with different clusters of users, with users within a given cluster having more homogenous preferences for many (app-category, permission) pairs. For instance, users in Cluster 1 are more permissive than users in Cluster 4.

[0083] The permission settings and/or privacy settings are intended to capture "user-specific preferences." They are needed when one-size-fits-all configurations would not adequately capture people's privacy preferences, for instance because not everyone in a given population feels comfortable granting a given permission or, more generally, configuring a given privacy setting in the same way. In some embodiments, the invention focuses on privacy preferences that indicate to what extent a user feels comfortable granting a given permission in a particular context. While many environments still lack such settings (e.g. at the time of writing, few if any camera monitoring systems allow users to opt-in to facial recognition, home assistants such as Amazon Echo or Google Chrome don't offer visitors the ability to opt out of having their voice recorded when they visit homes where these devices are deployed), there is a trend towards making such settings available to users, as people have generally been shown to have diverse privacy preferences. In other words, one-size-fits-all default settings are often unlikely to satisfy all users,

leading to either overly protective configurations, where a subset of users would have been willing to grant some permissions but the permissions are denied, or overly loose configurations where permissions are granted while some users do not feel comfortable with such granting. This trend is best illustrated by the evolution of mobile app permissions in operating systems such as iOS and Android. It is also visible in browsers and on different social networking sites, which have seen an increase in the number of privacy settings exposed to users. In addition, recent privacy regulations (e.g., CCPA, GDPR) have introduced requirements for giving users different permissions under different conditions (e.g., opting out of the “sale” of one's data under CCPA, or obtaining explicit opt-in for some practices under GDPR). In some embodiments, the invention focuses on helping users configure what often amounts to an unmanageably large collection of such privacy settings. The invention focuses on learning people's individual privacy preferences based on one or more sources of information, including information about the preferences users express towards a small number of representative data collection and use scenarios, information that can be used to compare a given user against others to identify groups of like-minded users and leverage their preferences, dialogues with users, settings the user has already configured, information about changes the user has made to some of his or her permission settings as well as other relevant historical data and other sources of information (e.g. using information about a user's privacy settings on Facebook or browser privacy settings to extrapolate his or her privacy preferences in other environments and help configure the corresponding privacy settings in these other environments).

[0084] In other embodiments, the modeling of privacy preferences is extended to also include preferences users have to be notified about different privacy practices that pertain to the collection and use of their data in different contexts. This is because often not all users want to be notified about the same data collection and use practices. For instance, one user might want to be notified every time he or she enters a room with camera monitoring systems that use facial recognition, while another user might only want to be notified about the presence of such camera monitoring system only in some particular environments (e.g. bars). One user might only want to be notified once about the presence of such camera monitoring system whereas another might want to be notified each time he or she is within the field of view of such a system. When privacy preference models are extended to include notification preferences and models of people's expectations (e.g. whether a user expects facial recognition to be used in conjunction with footage captured by a camera monitoring system, or whether a user expects his location information to be shared with third parties, or his conversation to be recorded), they can be used by Privacy Assistants to selectively determine when to notify their users about different types of practices. Models about how to notify different people can also be used (e.g. some users might prefer their smartphone to vibrate and display messages; others might prefer a sound and information to be displayed on their smart watch, etc.). As should be clear from the description, the PPA can interface with the user using any suitable input modality. For example, the computing device could have a display screen and the PPA could receive the user's inputs via text or writing; and/or the computing device could have a local or cloud-based speech recognition system that recognizes the user's utterances.

[0085] In some embodiments, the PPA **230** will include, or be configured to display, a dashboard that allows users to review and adjust their privacy settings. Examples of such a dashboard are shown in FIGS. **20A-20D** and described further below. In some embodiments, the dashboard may further allow users to review actions taken based on these settings such as the number of times some data practices were allowed, denied or restricted over a time period based on the user's current privacy settings (e.g., the number of times collection of the user's location was denied or allowed for one or more mobile apps based on the user's privacy settings). This could also be in the form of the number of times collection of a user's information by one or more IoT resources was allowed or denied as the user came within range of that IoT resource. For instance, a user might have a privacy preference to automatically opt-out of the collection of video footage about them

when such collection is for marketing purposes. Based on this preference, the user's PPA might have submitted opt-out requests on behalf of the user to five different IoT resources engaging in this practice over the past twenty-four hours, or over some other period of time (e.g., previous week). Some privacy settings shown on the dashboard might be specific to a given resource (e.g., a given mobile app or a given IoT resource such as a particular video analytics system in a particular store), others may apply to collections of such resources (e.g., all social media apps, or all video analytics resources deployed at stores belonging to a particular chain, or all Ring door bells, independently of their location). Yet other entries in the dashboard might apply to entire categories of privacy settings (e.g., opting out of the sharing of any personally identifiable information for marketing purposes independently of the type of personally identifiable information or the type of resource collecting this information). In some embodiments, the PPA's dashboard might allow the user to select among several different possible views such as only displaying certain types of resources, certain types of privacy settings, specifying a particular lookback horizon, and more. Some of the privacy settings displayed in the dashboard might be privacy settings directly entered by the user. Others may be privacy settings inferred using machine learning with the user potentially having the option to review and edit these settings before they are used by the PPA. In some embodiments, the PPA's dashboard may be available on the user's computing device such as in the form of screens in a mobile app implementing the PPA. In other embodiments, the dashboard may be accessible through a web browser. Under these different embodiments, the user is typically assumed to have been authenticated prior to being able to access their dashboard.

[0086] In some embodiments, for privacy reasons, information provided through the dashboard such as the number of times different data practices were allowed or denied over a given period of time, including possibly the specific IoT resources for which these practices were allowed or denied and possibly the location of these different IoT resources, will be displayed based on historical data collected and stored locally on the user's computing device (e.g., mobile phone, smart watch or other personal device carried or worn by the user). In some embodiment, the dashboard might include additional information such as the status of different privacy requests submitted by the user's PPA assistant to different IoT resources (e.g., status of deletion requests, status of opt-out requests) including status information such as “deletion requested but not yet confirmed” or “opt-out request submitted; IoT resource waiting for supporting documentation.” The latter example could correspond to a situation where the IoT resource is requiring proof that someone is a California resident before accepting a CCPA “do not sell my data” opt out request.

[0087] In some embodiments, the PPA dashboard may also send notifications to the user in the form of summaries. These summaries could include information such as what privacy settings or requests the PPA has communicated to different types of resources, how many resources of a given type the user's PPA has interacted with over a given lookback period, and other relevant data. These summary notifications may be implemented to act as nudges, prompting users to access their dashboard, review additional information and possibly modify their current privacy preferences (e.g. a user realizing that their footage has been captured by a number of video analytics systems for non-surveillance purposes and deciding to opt out of this type of collection when an opt out setting is made available by video analytics resources). The notifications could, for example, include notifications on a mobile device on which the PPA is running or notifications on other devices (e.g., smart watch or other wearable technology). The notifications could include and/or combine textual, visual, audio and haptic elements.

[0088] In various embodiments, where the PPA recommends possible privacy preferences to the user (e.g., based on machine learning models), the user could review, accept, modify or reject privacy preference recommendations for individual resources or groups of resources through their dashboard.

[0089] In addition, in various embodiments, the privacy preferences (and hence also the privacy recommendations) may not be limited to “allowing” or “denying” some data practices, but could

also include modulating the resolution at which data can be collected or the frequency at which some data is collected, or limiting the purposes for which collected data can be used or entities with which it can be shared. For example, in the case of mobile app privacy permissions, instead of allowing access to the user's fine location (e.g., based on a GPS unit on the user's computing device), a privacy recommendation for an app might be that the user's location only gets collected at a reduced level of resolution (e.g., no less than 100 yards, etc.). Similar reductions in resolution could apply to pictures of the user (captured by a camera of the computing device), the quality of sound captured by the computing device, and/or the frequency at which information is collected (e.g., location). In some embodiments, functionality may allow users to select settings that add noise or even entirely modify the user's data (e.g., falsifying the user's location, changing the computing device's ID, or transmitting intentionally erroneous data about the user). Similar recommendations for privacy preferences that restrict the collection and/or use of data by IoT resources can be supported in the context of a PPA assisting users manage user-specific privacy settings of IoT resources that may collect data about them.

[0090] In addition, privacy preferences used by the PPA to select which privacy setting to communicate to different IoT resources could be contingent on other factors or contextual attributes, such as the time of day and/or user's location. For example, a contingent preference (and hence also a corresponding preference recommendation) could be that an IoT resource is only allowed to collect some data about the user or engage in some other practice using that data during certain time periods and/or when the user is at certain locations (or not at certain locations), or conversely that engaging in some data practices is not allowed under specific conditions. Preferences for similar restrictions could be captured and communicated by a mobile app permission PPA. For example, access to the user's location information may only be granted to some apps when the user is not at home or at work, and/or at certain times of day or days of the week. Possible contextual attributes for contingent privacy preferences for an app, therefore, could be (i) a time of day attribute such that the app is granted or denied access only during certain times of the day; (ii) a day of week attribute such that the app is granted or denied permission to access some information or some functionality only during a certain day of the week; (iii) a location attribute such that the app is granted or denied permission to access information or functionality only when location information from the computing device indicates that the computing device is in a certain location or within a certain geographical area; (iv) an activity attribute such that the app is granted or denied permission when the user is engaged or not engaged in a particular activity (e.g., on a phone call, getting driving navigation instructions, etc.); and (v) an "in company of" attribute such that the app is only granted or denied permission when the user is in the company of or not in the company of certain people or categories of people (which itself could be determined based on location, calendar data, third party venue check-in data, or video analytics). In short, a number of different attributes can be used to capture people's often complex privacy preferences. Similar contingent permissions could be applied to restrict access to sensitive functionality on the user's computing device such as accessing the device's camera, a heart rate sensor, whether on the device itself or on some other device (e.g. smartwatch paired with the user's smartphone). Preferences and associated user-specific privacy settings for similar restrictions can also be modeled and submitted by a user's PPA to IoT resources (e.g., a video analytics resource only allowed to capture footage of the user when the user is not accompanied by a child).

[0091] A recommended privacy preference can also comprise, additionally or alternatively, at least one purpose attribute such that collection and use of the user's data is only allowed when the purpose attribute is satisfied. The purpose attribute could be things such as (i) sharing the information made accessible by the permission only with a specific category (or categories) of third parties; or (ii) a guarantee that the information made accessible by the permission will not be shared with specific category (or categories) of third parties.

[0092] The recommended permission setting can also comprise, additionally or alternatively, at

least one retention attribute such that the permission is only granted when the retention attribute is satisfied. The retention attribute could be, for example, a guarantee that information made accessible by the permission will not be stored for longer than a certain period of time.

[0093] In another embodiment, instead of determining privacy recommendation(s) based solely on the assignment of the user to a cluster (or other grouping), the recommendation(s) could be based on matching the user to a collection of clusters and by using privacy recommendation profiles (e.g., collections of recommendations) to identify recommended settings for the apps present on the user device as well as privacy settings associated with other technologies with which the user interacts. For instance, a user might be assigned to a cluster tied to his privacy settings on a social networking site or to questions related to his willingness to share different types of information in the context of social networking sites, and that same user may also have been matched with a cluster from a different collection of clusters related to his privacy preferences in the context of different Internet of Things scenarios. By combining preferences from these two clusters, it might be possible to infer a more complete and/or more accurate set of recommendations for that user. For instance, while profiles tied to each of these two clusters may not be sufficient to recommend a particular privacy setting (e.g., because data available for test subjects in each of these two clusters is not sufficiently conclusive to make a recommendation for the particular privacy setting), knowing that the user belongs to both clusters might provide sufficient support for making a recommendation for the setting. In other embodiments, additional recommendations are identified by aggregating new information collected from the user with information already collected from the user. Such additional information can include particular permission recommendations the user has accepted, rejected or modified over time, additional permissions the user has configured over time (possibly without the intervention of the personalized privacy assistant) and possibly answers to additional personalized questions asked over time by the privacy assistant.

[0094] As mentioned above, privacy preference models could be generated using a statistical analysis technique or a machine learning technique, such as a clustering technique. Examples of relevant machine learning and statistical analysis techniques that can be used to develop privacy preference models include: clustering techniques, collaborative filtering techniques, content-based filtering techniques, logistic regression techniques, support vector machine techniques, Bayesian inference techniques, decision tree learning techniques, and more, including ensemble methods that combine multiple models. Some elements of the preference models can also be improved using deep learning techniques. Other machine learning techniques such as unsupervised learning techniques can also contribute to further enriching collective privacy preference models. Other variations of machine learning techniques that can be used to build both collective and individual privacy preference models are disclosed in B. Liu et al., "Reconciling Mobile App Privacy and Usability on Smartphones: Could User Privacy Profiles Help?," 23rd International Conference on the World Wide Web (WWW 2014), July 2014, which is incorporated herein by reference.

[0095] In another embodiment, the data center could comprise an app performance database **113** (see FIGS. 2 and 11). This database may include data about how various apps perform when the permissions or settings are changed, such as whether the apps crash or otherwise lose functionality (e.g. a mobile comparison shopping app may continue to work but may not be able scan bar codes anymore if it is denied access to camera functionality, or a restaurant recommendation app may continue to work but may need to request the user to enter his/her location if it is denied access to fine grained location functionality). The app performance data can be collected from any suitable source, such as analysis of the app's code, test data from the app (e.g., running the app and monitoring its behavior, with and without certain permissions), information provided by the app developer (e.g. the app developer might indicate which permissions are absolutely required and/or what happens when different permissions are denied), and/or user app reviews, which may include comments about what happens when the app is denied a particular permission or recommendations to deny some permissions. As such, the privacy recommendations from the remote server **108** of

the data center **106** could be based on both the user's privacy preference model and the app performance data in the app performance database. That is, in one embodiment, the remote server **108** does not make a privacy recommendation if the recommendation would result in the app on the user's computing device crashing or otherwise losing critical functionality, as determined from the data in the app performance database. In that situation, the recommended permission settings might only include recommendations that have not been identified as permission settings that lead to operational problems in the corresponding app, where the operational problems could be causing the app to crash, causing the app to drain battery life at an unacceptable rate, or causing loss or unacceptable degradation of critical functionality in at least one of (i) the app itself, (ii) the computing device, and (iii) other apps installed on the computing device.

[0096] In another general aspect, the system is for notifying users about the presence of IoT resources (e.g. IoT devices or IoT systems) that collect information about the user and for optionally configuring user-specific privacy settings for these IoT resources, if such settings are available (see FIG. 1B). In some embodiments, the system comprises a Personalized Privacy Assistant (PPA) running on the user's computing device. The PPA is able to discover IoT resources listed in registries that are relevant to the user's current location (e.g. by allowing the PPA to submit to the registries queries requesting lists of IoT resources within the vicinity of the user's current location). For each of these IoT resources, the PPA is also able to query the resource's individual entry in the registry. Such entry typically includes a description of the resource such as who its owner is, what data it collects, how long it retains data and in what form (e.g. anonymized or not), how the data is used, who it is shared with, and whether the resource offers any specific APIs that enable individual users (via their PPA) to configure any available user-specific settings (e.g. each user having access to opt-in, opt-out privacy settings pertaining to one or more data collection and/or use practices). Based on these descriptions and based on models of the user's privacy notification preferences, the PPA can selectively determine whether the presence of the resource and its data collection and use practices warrant notifying the user. Privacy preference models may also be maintained by the PPA to help determine how to notify the user (e.g. different notification preferences such as different preferences for when to be interrupted and/or for one particular notification format over another). For instance, a privacy notification preference model may indicate that some users always want to be notified about the presence of video cameras that use facial recognition, or that some users do not care to be bothered by such notifications unless some other conditions apply (e.g. user being with a particular individual, user being at a sensitive location such as a bar, or when it is the evening, or when data is stored for more than a week). Privacy preference models can also be used to help the user configure any user-specific privacy settings exposed by a given resource (e.g. automatically opting out of some practices on behalf of the user, or recommending to the user that they opt in or opt out of various data collection and/or use practices). As already discussed earlier, privacy preference models, whether to determine when to notify users, how to notify them and about what practices as well as privacy preference models to help users configure user-specific privacy settings can be stored in whole or in part in a data center or on the user's computing device. As already indicated earlier, in general, these models can optionally be developed and refined through a number of different machine learning and/or statistical analysis techniques, which may also involve different forms of dialogues with users (e.g. asking users some questions, requesting users to confirm, clarify, refine, or generalize some of their decisions, nudging users to re-examine prior decisions, etc.). FIGS. 24A and 24B display clustering results showing how users can be organized in like-minded clusters when it comes to whether or not they want to allow or deny different types of video analytics practices. These clusters in turn can be used by a PPA to recommend privacy preferences to their users, with these preferences driving the submission of privacy requests to relevant IoT resources—in this case video analytics resources. As the PPA helps its user configure user-specific privacy settings exposed by different IoT resources and advertised in IoT resource registries, it accesses user-specific privacy setting

APIs advertised in the IoT resource registries for the corresponding IoT resource.

[0097] In some embodiments, these APIs support functionality to identify available privacy settings and more generally privacy requests supported for a given IoT resource- or possibly a collection of resources. In some embodiments, the APIs support functionality to check the current configuration of privacy settings for a given user as well as the status of privacy requests that a particular user may already have submitted to that IoT resource (e.g., whether a request for a copy of one's data or a request to have one's data deleted has been processed or whether supporting documentation might be required to verify that this user is entitled to submitting such a request, status of an opt-out request, etc.). These APIs typically also support functionality to modify the configuration of a privacy setting (e.g. turning off an opt-out setting) or modify a privacy request (e.g., canceling a request for deletion of one's data while the request has not yet been processed). In some embodiments, these APIs support submission of supporting documentation required for a privacy setting to be modified or more generally a privacy request to be honored such as proving that one is a resident of California and therefore qualifies to be treated under California Consumer Protection Act (CCPA), e.g., submission of CCPA opt-out or deletion requests.

[0098] In various implementations, as indicated above, the discovery of the third party IoT resources by the personal privacy assistant is mediated by registries that can be discovered and queried based on the location of the user (e.g. registry for a mall, for an office building, for a neighborhood, for a common collection of IoT resources, for IoT resource entries vetted or curated by a given individual or organization). Discovery could also be filtered by the user's PPA (e.g., based on the types of IoT resources the user is interested in, the particular data practices they want to be informed about, the availability of privacy settings, or the particular registry or registries in which some IoT resource descriptions are published). The entries in the registries can include descriptions of user-specific privacy settings and more generally privacy requests made available by a given IoT resource (e.g. opting in, opting out, modulating the granularity at which data is collected and/or stored, limiting the retention of some data, limiting the sharing of data with third parties, limiting potential use of the data, requesting that some data be erased, requesting a copy of the data collected by the resource, and more). These descriptions of user-specific privacy settings associated with an IoT resource can include a description of interfaces (or APIs) available for the personal privacy assistant to access and/or configure these settings and or more generally to submit user-specific privacy requests (e.g., not just requests involving opt-in or opt-out settings, but also requests for the deletion of one's data, requests to review data collected about oneself, requests for a copy of one's data). In some embodiments, user-specific permissions for third party IoT resources are maintained and enforced directly by the IoT resource. In other embodiments, they may be maintained and enforced by external Policy Enforcement functionality. Such functionality may include functionality to keep track of settings selected by different users and may also include functionality responsible for enforcing these settings. In some embodiments, Policy Enforcement functionality may be split into multiple components, possibly operating on different computer systems with these systems possibly provided by different entities (e.g., different organizations). For instance, a cloud-based consent management platform might be used to capture privacy settings and more generally privacy requests submitted by a user or their privacy assistant, whereas a separate module might be responsible for ensuring that the collection and use of a user's data is performed in accordance with the privacy settings communicated to the consent management platform.

[0099] In some embodiments, the IoT Privacy Infrastructure **300** provides a portal, which may support a variety of different roles, from system administrators for the entire infrastructure, IoT registry owners/administrators, IoT resource owners, administrators and contributors, IoT resource template owners, and PPA users, each provided with different views, different elements of functionality and some support to participate in different workflows associated with the portal (e.g., submission of IoT resource descriptions for review and publication, request for the creation of an

IoT resource registry, etc.). In some embodiments, the IoT portal allows resource owners (e.g., owners of IoT resources, mobile app publishers, etc.) and administrators (including specific employees of organizations owning or managing these resources) to publicize the data practices of the IoT resources. In some embodiments, owners or administrators of IoT resources access the portal to describe their IoT resources and their data practices, including descriptions of APIs available to support the discovery of these resources and their data practices, and submit requests to configure privacy settings for these resources or more generally submit privacy requests for these resources. In some embodiments, the discovery of resources and their data practices (including the privacy settings and more generally privacy requests they support) may rely in part or in whole on information directly stored in databases hosted within the IoT Infrastructure. In other embodiments, information stored in the description of an IoT resource in the database may include details on how to access external APIs (e.g., a URL and possibly a description of services hosted at that URL), which can be accessed to obtain additional information about a given IoT resource (e.g., the specific privacy settings and more generally privacy requests it supports, default configurations of privacy settings, documentation required to be able to take advantage of some settings or some requests). In some embodiments, the portal may allow users to request the creation of registries for collections of resource descriptions they want to manage, curate and/or publish (e.g., collection of resources in a mall, in a city, in an office building, in a neighborhood, in a private home, in one's car, etc. but possibly also collections of resources vetted by a given organization or collections of resources that share some other common attributes). The portal may support functionality that enables a system administrator to review such requests, including the geographic area possibly associated with one such request (i.e. creating a registry limited to the publication of resource descriptions in a given area), and determine whether to grant the request, ask for more information from the requestor to determine whether to grant the request, deny the request, modify some elements of the request (e.g., restricting the geographic area covered by a given registry).

[0100] In some embodiments, a system administrator may have the right to review resource descriptions before these descriptions actually get published by the system. Similar functionality might also be available to individual registry owners or administrators, enabling them to vet requests to publish IoT resource descriptions in their registry. Different registry owners/administrators may follow different policies when it comes to deciding what is required for an IoT resource description to be published in their registry. Information about the particular registry in which an IoT resource is published might be made available to privacy assistants, which in turn may follow different policies themselves such as different policies to decide which IoT resource and which practices to notify their users about (e.g., based on whether their users deem specific registries and their administrators to be sufficiently trustworthy, or simply based on the categories of resources published in specific registries or some other relevant criteria). A user might configure their privacy assistant **230** to only discover and interact with IoT resources in registries they consider to be sufficiently trustworthy or they deem to be relevant (e.g., a user deciding to only pay attention to IoT resource descriptions published in an official city registry and ignoring other descriptions not published in that registry). In some embodiments, the portal may possibly allow people who do not control or own IoT resources to report on the presence of resources and include partial or complete descriptions of these resources (e.g., someone noticing a camera system in a mall, but not knowing exactly what happens to the footage taken by the camera such as the purpose for which it is collected, how it is processed, for how long it is retained, with whom it might be shared). In some embodiments, the portal will want to make sure that everyone who enters IoT resource descriptions and requests to publish these descriptions has been properly vetted and authenticated. In some embodiments such an approach might be supplemented with ratings and system reviews of people who contribute descriptions of IoT resources they do not own or administer. In some embodiments, publication of a resource description, whether by the owner/administrator of that resource or by a third party contributor, might require the publisher,

whether owner/administrator or third party contributor, to certify that the description they provide is accurate to the best of their knowledge.

[0101] In some embodiments, the portal used by people to provide resource descriptions might also allow for the creation of resource templates, namely pre-filled descriptions of resources, including their data practices, APIs to communicate privacy settings (or more generally privacy requests) and other relevant information. The templates could then be used to quickly describe instances of these resources by just completing the missing information such as the particular location where a given instance of the resource template is being deployed. These templates may be created by a given organization solely for its own use (e.g., a system integrator that deploys similar systems at different locations and wants to save the effort of entering the same information over and over again). Templates could also be provided by device manufacturers (e.g. a manufacturer of a popular video door bell, a popular smart speaker, a particular car manufacturer, etc.). These templates could then be used by people who purchase and deploy these resources, allowing them to quickly instantiate the template to reflect the particular location where the resource is deployed and the particular configuration of that resource such as the particular settings (or more generally privacy requests) made available to people whose data might be collected by the resource and their default values (e.g., whether a video door bell has been configured to share its footage with law enforcement or not). By enabling users to instantiate these templates, it is possible to significantly reduce the amount of effort required to publish the presence of these resources. In particular, most people do not have the time or expertise required to fully understand the data practices of IoT resources they purchase, which would typically require reading the privacy policy of these resources. By providing a template that captures these practices and makes them ready for publication, the manufacturer not only alleviates the burden of publishing the presence of its IoT resources and their data practices but also significantly reduces the chances that the IoT resource description is inaccurate.

[0102] The components in the different embodiments discussed herein should be implemented to meet relevant security requirements and communication between components should also meet relevant security properties (e.g., authentication, authorization, confidentiality, integrity, non-repudiation, availability, etc.). Different protocols may be used to meet such requirements, with some solutions offering stronger guarantees than others. It should be noted that different contexts often entail different security requirements. For instance, simply discovering IoT resources near one's location and their data practices may not necessarily require that users be authenticated, though under some scenarios it might (e.g., scenarios where the presence of some resources and their data practices should only be communicated to people authorized to know about their presence). In some embodiments, a user could rely on a privacy assistant app (e.g. a privacy assistant app downloaded from an app store, such as shown in FIG. 21) to discover IoT resources around them and the data practices associated with these resources without being required to authenticate with the privacy assistant. The privacy assistant itself may or may not be required to authenticate and may or may not need to be authorized to submit requests for information to registries or to IoT resources, though in general it is safer to require some level of authentication and/or authorization (e.g., to mitigate potential denial of service attacks).

[0103] On the other hand, when a user or a privacy assistant acting on behalf of a user submits a user-specific permission setting or more generally a privacy request to a third party IoT resource or the entity controlling that IoT resource (e.g., request to opt in or out of a particular data practice, request to have one's data deleted, request to review one's data), the IoT resource or the controlling entity may require proof that the user and/or the privacy assistant are indeed who they claim to be and/or proof that they are allowed to submit the particular user-specific request or privacy setting. A number of different conditions for authentication and/or authorization and a number of implementations may be used according to embodiments of the present invention.

[0104] In some embodiments a web server (or a collection of servers) acts as an overall "privacy

infrastructure” system, hosting one or more registries, IoT resource descriptions, IoT resource templates, and also taking care of authenticating different users (e.g., registry owners/administrators, IoT resource owners/administrators as well as contributors, PPA users, overall system administrator(s), IoT resource template owners, etc.)

[0105] In some embodiments, when receiving a privacy request from a user or some other communication from a user, an IoT resource or its controlling entity will want to verify the identity of the user and/or verify that the user meets some conditions (e.g., being a resident of a particular State to be allowed to submit a CCPA privacy request, being in a particular mall, being over a certain age). In some contexts verification may extend to obtaining approval from a third party and verifying such approval (e.g., approval from a parent or guardian). In some situations, such verification may be limited to just authenticating the user, which may be done by the privacy infrastructure. In other situations, such verification may also require the submission of additional digital data (e.g., electronic documents proving one's age or residence, sensor data such as facial recognition data, location data if submission of some privacy settings is limited to people present in a given area such as a room or building, etc.). This digital data may be in the form of signed certificates or possibly through copies of documents (e.g., copy of one's driver's license or utility bill). Submission of this digital data could be handled by the privacy assistant on behalf of the user if such data is accessible by the privacy assistant, or it may have to be done by the user, possibly through external communication channels. In general each entity involved in the submission of privacy requests and of supporting digital documentation required for a given request will need to be authenticated. These interactions may be in part or in whole mediated by an authorized agent, which may be part of the privacy infrastructure or may be a third party authorized agent. Different IoT resources may require working with different authorized agents. A list of one or more authorized agents recognized by an IoT resource or the entity controlling an IoT resource may be included in the description of the IoT resource as stored in the IoT infrastructure. Upon discovering an IoT resource, a privacy assistant would determine which authorized agent the IoT resource is able to work with and would submit any user-specific privacy request to one of the authorized agents recognized by the IoT resource. In some instances, the IoT privacy infrastructure might itself be recognized as an authorized agent by an IoT resource and might then be able to directly forward user-specific requests submitted by the user's privacy assistant to the corresponding IoT resource. In some embodiments supporting documentation required for some user-specific privacy requests may be stored by a user's privacy assistant and submitted by the privacy assistant directly to an IoT resource or indirectly via an authorized agent. In some embodiments some supporting documentation might be directly accessible by the authorized agent, which could then directly submit it to the IoT resource.

[0106] The privacy assistant may store and/or facilitate the transmission of digital data necessary to recognize the user, verify that the user is entitled to submitting a specific privacy/preference setting or request (e.g., opting in or out of a data practice, requesting some restrictions over a particular data practice, requesting deletion of one's data, requesting to review one's data stored by an IoT resource) and help ensure that an IoT resource's data practices implement the user-specific permission setting or privacy request (e.g., by waiting for an acknowledgement from the IoT resource and/or handling follow-on communication with the IoT resource such as requests for additional information.). In some embodiments, the privacy assistant might store digital credentials necessary to recognize the user (e.g., credentials to demonstrate residence in a given jurisdiction, credentials associated with membership in a given organization, credentials showing that the user has been adequately authenticated, credentials confirming that the user meets a minimum age requirement).

[0107] As already indicated authorized agents may help mediate interactions between a user's privacy assistant and an IoT resource (or an entity controlling the IoT resource). The privacy infrastructure itself may act as an authorized agent recognized by a number of IoT resources but

some IoT resources may require working with third party authorized agents. When the IoT infrastructure is recognized as an authorized agent, there is no need for the user to further authenticate, as the user will already have authenticated with the IoT infrastructure. In some other embodiments, the IoT Infrastructure might already be trusted by the authorized agent. In such situations, the authorized agent may be able to rely on the fact that the user has already been authenticated by the IoT privacy infrastructure, which in turn helps streamline the process of submitting privacy requests on behalf of the user to IoT resources (assuming that these resources trust the particular authorized agent).

[0108] In yet other embodiments, the privacy assistant might store or help transmit digital information needed to recognize data collected about a particular user and ensure that this data is processed by the IoT resource in accordance with the user-specific permission setting submitted on behalf of that user. Such information could take many different forms. For instance, it could be a MAC address in the case of a Wi-Fi location tracking resource using MAC addresses to track people's whereabouts (by effectively tracking the location of their devices). In the case of a video analytics system used to recognize people it could be biometric data such as one or more photos of someone's face or footage of that person's face, with the privacy assistant communicating this information to the IoT resource along with a privacy settings (e.g., a user opting in and allowing a video analytics resource to track their whereabouts in a store or track their facial expressions as they look at items behind store windows in a mall).

[0109] In some embodiments, such as shown in the example of FIG. 17, the privacy assistant **230** in the system may be implemented as a module that is separate from the authorized agent **310**, with both of these modules possibly operated by different parties, each intended to assist the users with a subset of the steps outlined above (e.g., the privacy assistant helping the user select relevant privacy settings, then communicating these settings to a separate authorized agent **310** for transmission to the relevant IoT resource **115** or an entity controlling that IoT resource). There are a number of different ways these authorized agents can be implemented and used to support the communication of user-specific settings and privacy requests to IoT resources or entities controlling these resources.

[0110] The example of FIG. 17 shows the PPA **230** on the user computing device **114** communicating via communication link **1001** with an IoT Privacy Infrastructure **300**, which can include the IoT resource registr(ies) **200** described above, IoT Resource Descriptions, IoT authentication services, and an IoT Portal service. The IoT privacy infrastructure **300** may also include, in various embodiments, an authorized agent, such as the authorized agent **310**. The IoT Privacy Infrastructure **300** may comprises servers (e.g., web server, application servers, database servers) and database(s). The IoT resource registry thereof might advertise the IoT resource **115** to nearby devices/PPAs as explained herein. The IT resource registry **200** IoT Privacy Infrastructure **300** may advertise via communication link **1001** the available user-specific privacy requests that can be made with respect to the data collected about the user by the IoT resource **115**. Those available user-specific privacy requests could comprise, opt-in and/or opt-out of data practices of the IoT resource **115**, deletion of the data collected about the user, and/or a request to review the data collected about the user as well as a number of other requests to restrict the collection and use of the user's data. Those available user-specific privacy requests could comprise, opt-in and/or opt-out of data practices of the IoT resource **115**, deletion of the data collected about the user, and/or a request to review the data collected about the user.

[0111] The IT resource description available to the PPA **230** via the IoT Infrastructure **300** may also advertise to the PPA **230** the authorized agent(s) **310** that are trusted by the IoT resource **115**. The PPA **230** may then request, via communication link **1002**, information about the data practices of the IoT resource **115** from IoT Privacy Infrastructure **300** and/or the authorized agent **310**, such as the available user-specific privacy requests that can be made with respect to the data collected about the user by the IoT resource **115**. The available user-specific privacy requests may be

transmitted by IoT Infrastructure via the communication link **1001** or by the authorized agent **310** to the PPA **230** at communication link **1002**. The authorized agent **310** may obtain the available privacy request from the IoT resource **115** at communication link **1003** in response to (i) a request at communication link **1002** from the PPA **230** and then (ii) a request from the authorized agent **310** to the IoT resource **115** for the latest, up-to-date available user-specific privacy requests of the IoT resource **115**. The authorized agent **310** may also query the IoT Privacy Infrastructure **300** for the latest, up-to-date available user-specific privacy requests of the IoT resource **115**.

[0112] Once the PPA **230** has the available user-specific privacy requests, the PPA **230**, using its privacy preference model for the user, for example, can recommend a privacy request for the user for the IoT resource **115**. In the embodiment of the FIG. **17**, the PPA **230** transmits the privacy request to the authorized agent **310** via communication link **1002**. The authorized agent **310** can then transmit, via communication link **1003**, the user's privacy request to the IoT resource **115** for implementation by the IoT resource.

[0113] In some scenarios, the IoT resource **115** determines that it needs documentation, verification, or other information that demonstrates that the user is authorized to submit the privacy request. For example, a user may need to be a certain age to opt-in to certain data collection practices of the IoT resource. The IoT resource **115** then requests the additional information from the authorized agent **310** via communication link **1003**. If the authorized agent **310** already has the required information, such as if the user has already submitted documentation that the user is a certain age via, for example, a digital image of the user's driver's license, the authorized agent **310** can respond to the IoT resource **115** with the requested information via communication link **1003**. If the authorized agent **310** does not already have the information, the authorized agent **310** can request it from the user via communications with the PPA **230** over the communication link **1002**. In response, the user can transmit, via the communication link **1002**, the requested electronic documentation (e.g., an image of the user's driver's license) to the authorized agent **310**, which in turn can transmit the electronic documentation to the IoT resource **115** via the communication link **1003** so that the user's privacy request can be implemented with respect to the data collected by the IoT resource **115**. In other scenarios, the authorized agent **310** can confirm that the user meets the criteria for the privacy request, such that the authorized agent **310** sends confirmation of approval to the IoT resource **105**, rather than the electronic documentation. The IoT resource **115** preferably has resource to process the electronic documentation sent to it to verify the privacy requests, such as machine learning classifiers, document readers, etc.

[0114] In some instances, information may be needed from a third party to authenticate the privacy request. For example, a minor may need approval from a parent or guardian to opt-in to certain data collection practices. As another example, to confirm a user's residence, the authorized agent **310** could query a third party system **311**, such as a utility, to confirm the user's address. In such circumstances, the authorized agent **310** can query a database for an address for the third party system **311** associated with the user/PPA and for the information requested. Once the authorized agent obtains the address, the authorized agent **310** and the third party system **311** can communicate via the communication link **1004** to obtain the necessary information to validate the user's privacy request. Then the authorized agent **311** can forward the third party information to the IoT resource via the communication link **1003**. In other embodiments some of this information might be provided by the privacy assistant or might possibly be submitted via external communication channel (e.g., email, uploading documentation to a secure portal, etc.).

[0115] In some instances, the documentation sent from the authorized agent **310** may include signed authorization token signed by the authorized agent **310**. The digital signature may be created using a private key held securely by the authorized agent. The token may be encoded and signed using standards like JWT (JSON Web Token). The IoT resource **115** can verify the digital signature using a public key of the authorized agent **310**. In other instances, the documentation may be digitally signed by the issuer of the documentation. For example, to prove that the user resides at a

certain address (in a certain zip code, county, states, etc.), an invoice from a utility company may be used. In such instances, the user may obtain a electronic copy of the invoice with a digital signature of the utility company, and such documentation could be used to prove that the user is authorized to submit the privacy request.

[0116] The authorized agent **310** may comprise one or more servers, such as a web server(s), an application server(s), and a database server(s). The authorized agent **310** could be implemented in a cloud computing environment, with the communications to and from it over secure channels (e.g., HTTPS with TLS). The authorized agent **310** could be part of the IoT privacy infrastructure **300** or it could be separate from (outside of) the IoT privacy infrastructure **300**. In various embodiments, the authorized agent **310** could also be integrated into the PPA **230** or other run on the user computing device **114**. The authorized agent **310** can act as an intermediary between the PPAs and IoT resources, particularly where an authorized agent is a trusted source of privacy request for an IoT resource(s). FIG. **17** only shows a single authorized agent **310**, but it should be recognized that multiple authorized agents may act as intermediaries for a single IoT resource. Also, each authorized agent could act as an intermediary for one or more IoT resources. The authorized agents **310** could also store verification (or electronic documentation verifying) different facts, such as a user's age, place of residence, etc., which verification could then be re-used for a subsequent privacy request for the user. Different documents may have different digital signatures and the authorized agent itself **310** could certify some of the documentation required to validate a privacy request. In that sense, the authorized agent **310** may provide decentralized trust management. The authorized agent may use immutable (e.g., blockchain) ledgers to record trust relationships and credentials (e.g., privacy request verifying information).

[0117] The IoT resource **115** may transmit timestamped electronic communications to the PPA **230**, e.g., via the authorized agent **310**, confirming that the user's privacy request was received, authenticated/validated, and/or implemented. Also, the above description of FIG. **17** assumed that the IoT resource **115** controlled and implemented the data practices of the IoT resource **115**. In other embodiments, a consent management platform or PEP associated with the IoT resource **115** may control and/or implement the privacy requests for the IoT resource, in which case the privacy requests and the authentication information can be sent to the consent management platform or PEP for the IoT resource **115**, as the case may be. A consent management platform (CMP) is a system that helps organizations collect, manage, and document user consent for data collection and processing. A CMP and a PEP can work together for an IoT resource. The user's privacy request can be submitted to the CMP, which records the privacy request. The PEP for the IoT resource checks the user's privacy request stored in the CMP and enforces the privacy request.

[0118] The information required to show that the PPA is authorized to submit the user request may comprise, for example, electronic documentation sufficient to show the required information about the user, such as age, residence in a particular jurisdiction, membership in an organization, etc. The information could be electronic files, such as digital images of driver's licenses or documents, pdf files of documents, digitally signed documents, etc. The IoT resource **105** may comprise, for example, machine learning systems trained to read the information from the submitted documents. The available user-specific privacy requests could be, for example, opt-in and opt-out request for data practices of the IoT resource (e.g., collection, retention, anonymization, identification practices, etc.), deletion of the data collected about the user, review the data collected about the user, etc.

[0119] The IoT Privacy Infrastructure **300** can comprise a web service that hosts a web page with a user interface for an administrator or owner of an IoT resource **115**, where the user interface provides an IoT resource template that is instantiated by a user of the IoT resource owner/admin computer **315** to publicize the presence of a particular IoT resource at a particular location. The template once instantiated (thereafter a "completed form") can include information about the data practices of the IoT resource as well as the various settings configurable by the user about those

data practices. Third parties can access the completed form online, via the internet, for information about the IoT resource, such as its data practices and privacy settings. For example, the IoT resource **115** could have a QR code near it or associated with it. A user could scan the QR code with his/her smartphone to access a web site that hosts the completed form. Also, the IoT resources may show in a user interface of the PPA **230** (such as based on vicinity or a user search), such as on a map, and the user could select one of the IoT resources in the app to be displayed the publicized, template-generated information about IoT resource.

[0120] In yet other embodiments, the present invention is directed to a system, comprising a web server that hosts a web page, or an application server that provides a dashboard, with a user interface for an administrator or owner of an IoT resource **115**, where the user interface provides an IoT resource template that is instantiated by a user from an IoT resource owner/admin computer device to publicize the presence of a particular IoT resource at a particular location. The template once instantiated (thereafter a “completed form”) can include information about the data practices of the IoT resource as well as the various settings configurable by the user about those data practices. Third parties can access the completed form online, via the internet, for information about the IoT resource, such as its data practices and privacy settings. For example, the IoT resource **115** could have a QR code near it or associated with it. A user could scan the QR code with his/her smartphone to access a web site that hosts the completed form. Also, the IoT resources may show in a user interface of the PPA **230** (such as based on vicinity or a user search), such as on a map, and the user could select one of the IoT resources in the app to be displayed the publicized, template-generated information about IoT resource.

[0121] In another general aspect, the present invention is directed to the IoT privacy infrastructure **300** that comprises an IoT resource registry, and web/application servers for providing the PPA and IoT resource portals described herein. The IoT privacy infrastructure **300** may also comprise one or more authorized agents.

[0122] FIGS. **18A-E** illustrate an example of the template that the owner/administrator could fill-in online to publicize the IoT resource. In the first page of the template, shown at FIG. **18A**, the owner/admin fills in basic information about the IoT resource. At the second page, shown in FIG. **18B**, the owner/admin can specify the location and geographic coverage of the IoT resource. At the third page, such as FIG. **18C**, the owner/admin can select from a predefined list types of data collected by the IoT resource, where the data collection is identifiable, and measure that will be taken before processing or publishing the collected data. On the fourth page, at FIG. **18D**, the owner/admin can specify the purpose for which the data are collected, how long the data will be retained, who will have access to the data, and the types of analytics that will be performed with the collected data. And on the fifth page of the template, shown at FIG. **18E**, the owner/admin can specify privacy options available to user, such as requesting data deletion, requesting a copy of the collected data, allowing or denying data collection, and/or allowing or denying sharing of the data. Once the owner/admin complete the template, the completed form can then be available online to users as described above. FIGS. **19A-E** show more example of the portal for the IoT resource administrators.—FIG. **19A** shows an example IoT Privacy Infrastructure dashboard for a system administrator who also manages five different IoT resource registries, has published an IoT resource description and has another IoT resource description that has not yet been published. This figure illustrates how a given user of the IoT privacy infrastructure **300** can have multiple roles, in this particular case, (a) overall administrator of the IoT privacy infrastructure, including the ability to review and grant request for the creation of IoT registries, (b) IoT resource registry owner, (c) IoT resource owner/administrator, and (d) IoT resource template owner/administrator, even though this user does not have any such template.

[0123] FIG. **19B** illustrates an example IoT privacy infrastructure portal showing a registry administrator a collection of IoT resources published in the registry. Each of the IoT resource descriptions can be selected to access all the details included in the description. The registry

administrator has access to functionality that enables them not just to review resource descriptions submitted for publication (in this particular case, all IoT resource descriptions were submitted by a contributor identified as “elouie”), but can also decide whether to accept or reject requests for publication of a resource description in their registry, and also unpublish a description at a later time, if warranted. The registry administrator has access to additional functionality, including transferring ownership of IoT resource descriptions to other users, transferring administration of the registry to others, and more. Workflow management functionality helps coordinate different roles on the portal, showing each user relevant status information in their dashboard (see FIG. 19A).

[0124] FIG. 19C shows another example view of a different IoT resource registry where the boundary of the registry can also be seen.

[0125] FIG. 19D illustrates an example view of IoT resource descriptions created by a given user of the IoT Privacy Infrastructure. This user has one resource that has been published and another that has not been published.

[0126] FIG. 19E shows an example view of an IoT resource description published in the IoT Privacy Infrastructure. Information about the URL and APIs to access privacy options has been masked for security reasons. The QR code is to allow resource owners/administrators to publicize the presence of the IoT resource to people who may not have the PPA mobile app on their smartphone. Someone scanning the QR code can first be brought to an app store (iOS or Google Playstore, depending on their mobile device), offered a chance to download the app, and, upon downloading the app, they are shown a description of this particular resource within the app, including access to the privacy options (i.e. privacy settings/privacy requests supported by the IoT resource, in this case the Lab Location Tracking resource). A user who already has the app and scans the QR code can be directly brought to a description of the IoT resource within the PPA app. This view is an example of IoT resource description and can include identification of the authorized agent(s) for the IoT resource.

[0127] FIGS. 22A-F illustrate additional portal screenshots for the IoT Privacy Infrastructure portal. FIG. 22A illustrates an example IoT Privacy Infrastructure portal showing a screen where an administrator can review requests for registries and their current status. The identity of requesters and requested names of some registries have been obfuscated for privacy reasons. FIG. 22B shows a sample screen of the IoT Privacy Infrastructure Portal showing a user submitting a request for an IoT Resource Registry in Pittsburgh. FIG. 22C shows a sample screen of an IoT Privacy Infrastructure Portal showing a system administrator reviewing a request by an authenticated user named John Smith for the creation of a resource registry in Pittsburgh. The system administrator can approve or decline the request. In the latter case, the system administrator can also provide feedback to the requester. FIG. 22D shows an example screen of an IoT Privacy Infrastructure Portal showing a resource owner who is about to pick a registry where to submit a request to publish their IoT resource description. The screen shows that the resource has already been published in a first register. It also allows users to self-publish, thereby optionally bypassing registries and taking full responsibility for the publication of the resource. FIG. 22E shows an example screen of an IoT Privacy Infrastructure Portal showing a prompt asking a resource owner to take responsibility for the content of the resource description and confirm that it is accurate to the best of their knowledge and compliant with applicable laws. FIG. 22F shows an example screen of an IoT Privacy Infrastructure Portal showing a prompt asking a resource owner to confirm upon the creation of a resource description to confirm that the description is accurate to the best of their knowledge and conforms with applicable laws.

[0128] FIGS. 20A-D show example screenshots of the PPA 230 according to various embodiments. The PPA 230 may be in communication with an application server of the IoT Privacy Infrastructure 300 to serve the information in the screenshots. FIG. 20A shows a map view of nearby IoT resources, color coded by their type of data category (e.g., visual, health, audio, etc.). The user can

select on one of the IoT resources in the map view to be shown more detail about the selected IoT resource, as shown in the example of FIGS. 20B-C, which show examples of the various information about the IoT resource that the user can see by scrolling through the description. And the example PPA view in FIG. 20D shows the status of the user's privacy requests for the selected IoT resource.

[0129] FIG. 23 shows an example screen of the PPA showing how a user can configure their notification settings, specifying the types of IoT resources (e.g. based on the type of data they collect, whether data is collected in a personally identifiable form, etc.) they want to be notified about as well as the frequency of these notifications (e.g., never, first time, every time).

[0130] FIGS. 24A and 24B illustrate a clustering analysis that can be used to organize users in groups of like-minded users. Privacy preferences associated with users in different clusters, such as whether or not to allow or reject different types of video analytics practices can be used by a PPA to recommend privacy preferences to a user based on the cluster that is the best match for them. These preferences in turn can then be used by the PPA to submit privacy requests to nearby IoT resources on behalf of the user, possibly directly or possibly with the help of an authorized agent. The example of FIGS. 24A and 24B show profiles associated with a 6-cluster model, namely a model that organizes a population of users into six groups of like-minded people with each group having a profile indicating whether people in this group are typically inclined to allow or reject different video analytics practices, with different practices being represented by different rows in the table. The third column for each cluster indicates whether people in this cluster are generally comfortable (dark grey, no dashed lines), uncomfortable (dark grey, horizontal dashed lines) or divided for each of the different video analytics practices. Each cluster profile contains three columns: the left one displays the average mean value (deny=-1, allow=1), and the right column represents the cluster profile, where a first color/hue (e.g. dark grey, no dashed lines) represents an allow decision, a second color/hue (e.g., dark grey, horizontal dashed lines) means a deny, and no color means no clear preference within a given cluster, either because there are not enough data points available or for lack of a clear majority. The middle column shows the variances, ranging from 0 to 1. The three numbers (D/A/T) in each entry in the right column represent the distribution of deny ("D") and allow ("A") collected for members of the cluster for the corresponding purpose, with T=D+A representing the total number of decisions collected for the given purpose from members of the cluster.

[0131] In one general aspect, therefore, the present invention is directed to a system: comprising an Internet of Things (IoT) resource for remotely sensing data about a user; and a computing device of the user. The computing device comprises a processor that executes a personal privacy app ("PPA"). The PPA is configured to: (i) receive data about the IoT resource, where the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; (ii) communicate a privacy request for the user with respect to the IoT resource, where the privacy request is one of the available user-specific privacy requests and where the privacy request communicated by the PPA is based on the data received about the IoT resource; and (iii) in response to a query related to the privacy request, causes electronic documentation to be transmitted that demonstrates that the user qualifies to submit the privacy request, such that the privacy request is applied to data collected about the user by the IoT resource.

[0132] In various implementations, data about the IoT resource could come entirely from the IoT resource description or may also include additional data obtained indirectly by using APIs provided in the IoT resource description and using these APIs to contact the IoT resource and obtain additional information about the resource's data practices, such as the particular privacy requests supported by the resource, the documentation required to demonstrate that the user qualifies to submit a given privacy request, and/or the default practices (e.g., privacy settings) of the IoT resource.

[0133] In various implementations, the PPA may communicate the privacy request to an authorized

agent, which can forward the privacy request for implementation with respect to the data collected about the user by the IoT resource. There can be one or more authorized agents available for a given IoT resource, and identification of those one or more authorized agents can be included in the IoT resource description. Also, the APIs to communicate with the one or more authorized agents can also be included in the IoT resource description. Thus, for instance, the PPA might discover an authorized agent for a given IoT resource, it might then contact that authorized agent to find what specific privacy requests are supported for that resource, it might then combine this information with the data practice information contained in the IoT resource description and integrate this information for display in the PPA to the user (or for the PPA to automatically communicate specific privacy requests on behalf of the user based on the user's privacy preferences).

[0134] In another general aspect, the inventive system comprises: the IoT resource for remotely sensing data about a user; a computing device of the user, where the computing device comprises a processor that executes a personal privacy app (“PPA”); and an authorized agent that is authorized to submit privacy requests on behalf of the user and is a trusted source of privacy requests for the IoT resource. The PPA is configured to receive data about the IoT resource, where the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; and communicate, to the authorized agent, a privacy request for the user with respect to the IoT resource, where the privacy request is based on the data received about the IoT resource. Also, the authorized agent is configured to communicate the privacy request for the user for implementation with respect to data collected about the user by the IoT resource.

[0135] Again, data about the IoT resource could come entirely from the IoT resource description or may also include additional data obtained indirectly by using APIs provided in the IoT resource description and using these APIs to contact the IoT resource and obtain additional information about the resource's data practices, such as the particular privacy requests supported by the resource, the documentation required to demonstrate that the user qualifies to submit a given privacy request, and/or the default practices (e.g., privacy settings) of the IoT resource.

[0136] In various implementations, in response to a query related to the privacy request, the PPA can cause electronic documentation to be transmitted to the authorized agent that demonstrates that the user qualifies to submit the privacy request. The authorized agent can then forward the documentation, such that the privacy request is applied to data collected about the user by the IoT resource.

[0137] In various implementations, the list of available privacy requests may be communicated to the PPA by the IoT privacy infrastructure **300**, such as the IoT resource registry **200** and/or the IoT resource description hosted by the IoT privacy infrastructure **300**.

[0138] In various implementations, the electronic documentation that shows that the user is qualified to submit the privacy request could comprise a digital image (e.g., of the user's driver's license, etc.) or an electronic document with a digital signature to verify authenticity of the document.

[0139] In various embodiments, the PPA is configured to determine the privacy request using a privacy preference model for the user.

[0140] In various implementations, the available user-specific privacy requests comprise opt-in and opt-out settings for data practices of the IoT resource, to request deletion of the data collected about the user, and/or to review (such as by being sent a copy of) the data collected about the user by the IoT resource.

[0141] In various implementations, the PPA also receives an electronic notification that the privacy setting was received, verified, and/or implemented with respect to the data collected about the user by the IoT resource.

[0142] In various embodiments, the system further comprises an IoT resource registry that advertises the IoT resource to the PPA. In which case, the PPA can discover the IoT resource registry based on location of the user computing device.

[0143] The servers **108** of the data center **106**, the IoT Privacy Infrastructure **300**, and the

authorized agent 310 may comprise one or more, preferably multi-core, processors and one or more memory units. The memory units may comprise software or instructions that are executed by the processor(s). The memory units that store the software/instructions that are executed by the processor may comprise primary computer memory, such as RAM or ROM, and/or secondary computer memory, such as hard disk drives and solid state drives. The software modules and other computer functions described herein may be implemented in computer software using any suitable computer programming language such as C#/.NET, C, C++, Python, Java, Javascript, Objective C, Ruby and using conventional, functional, or object-oriented techniques. Programming languages for computer software and other computer-implemented instructions may be translated into machine language by a compiler or an assembler before execution and/or may be translated directly at run time by an interpreter. Examples of assembly languages include ARM, MIPS, and x86; examples of high level languages include Ada, BASIC, C, C++, C#, COBOL, Fortran, Java, Lisp, Pascal, Object Pascal, Haskell, ML; and examples of scripting languages include Bourne script, JavaScript, Python, Ruby, Lua, PHP, and Perl.

[0144] The examples presented herein are intended to illustrate potential and specific implementations of the present invention. It can be appreciated that the examples are intended primarily for purposes of illustration of the invention for those skilled in the art. No particular aspect or aspects of the examples are necessarily intended to limit the scope of the present invention. Further, it is to be understood that the figures and descriptions of the present invention have been simplified to illustrate elements that are relevant for a clear understanding of the present invention, while eliminating, for purposes of clarity, other elements. While various embodiments have been described herein, it should be apparent that various modifications, alterations, and adaptations to those embodiments may occur to persons skilled in the art with attainment of at least some of the advantages. The disclosed embodiments are therefore intended to include all such modifications, alterations, and adaptations without departing from the scope of the embodiments as set forth herein.

Claims

1. A system comprising: an Internet of Things (IoT) resource for remotely sensing data about a user; and a computing device of the user, wherein the computing device comprises a processor that executes a personal privacy app ("PPA") that: receives data about the IoT resource, wherein the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; communicates a privacy request for the user with respect to the IoT resource, wherein the privacy request is one of the available user-specific privacy requests and wherein the privacy request communicated by the PPA is based on the data received about the IoT resource; and in response to a query related to the privacy request, causes electronic documentation to be transmitted that demonstrates that the user qualifies to submit the privacy request, such that the privacy request is applied to data collected about the user by the IoT resource.
2. The system of claim 1, wherein the PPA receives the data about the IoT resource from an IoT resource description for the IoT resource hosted by an IoT privacy infrastructure.
3. The system of claim 2, wherein the PPA receives the data about the IoT resource via API calls to the IoT resource, wherein an API for the API call is included in the IoT resource description.
4. The system of claim 1, wherein the data about the IoT resource comprises privacy requests supported by the IoT resource.
5. The system of claim 1, wherein the data about the IoT resource comprises identification of documentation that demonstrates qualification to submit the privacy request.
6. The system of claim 1, wherein the data about the IoT resource comprises default privacy practices of the IoT resource.
7. The system of claim 1, wherein the PPA submits the privacy request via an authorized agent,

wherein the authorized agent is a trusted source of privacy requests for the IoT resource.

8. The system of claim 7, where the PPA causes the electronic documentation to be transmitted via the authorized agent.

9. The system of claim 8, wherein the PPA receives the data about the IoT resource from an IoT resource description for the IoT resource hosted by an IoT privacy infrastructure, and wherein the IoT resource description: identifies one or more authorized agents that are trusted sources for privacy requests for the IoT resource; and APIs to communicate with the one or more authorized agents.

10. The system of claim 1, wherein the PPA is further configured to determine the privacy request for the user using a privacy preference model for the user and using the data about the IoT resource.

11. The system of claim 1, wherein the electronic documentation comprises a digital image.

12. The system of claim 1, wherein the electronic documentation comprises an electronic document digitally signed.

13. The system of claim 1, wherein the electronic document evidences an attribute of the user, wherein the attribute is one of an age of the user, a location of residence of the user, membership of the user in an organization, or a clean criminal record.

14. The system of claim 1, wherein the available user-specific privacy requests at least one of opt-in or opt-out settings for a data practice of the IoT resource, a request for deletion of the data collected about the user, or a request to review the data collected about the user.

15. The system of claim 1, further comprising an IoT resource registry for advertising the IoT resource, and wherein the PPA is for discovering the IoT resource registry based on a location of the computing device of the user.

16. A system comprising: an Internet of Things (IoT) resource for remotely sensing data about a user; a computing device of the user, wherein the computing device comprises a processor that executes a personal privacy app (“PPA”); and an authorized agent that is authorized to submit privacy requests on behalf of the user and is a trusted source of privacy requests for the IoT resource, wherein the PPA: receives data about the IoT resource, wherein the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; and communicates, to the authorized agent, a privacy request for the user with respect to the IoT resource, wherein the privacy request is based on the data received about the IoT resource, and wherein the authorized agent communicates the privacy request for the user for implementation with respect to data collected about the user by the IoT resource.

17. The system of claim 16, wherein the PPA receives the data about the IoT resource from an IoT resource description for the IoT resource hosted by an IoT privacy infrastructure.

18. The system of claim 17, wherein the PPA receives the data about the IoT resource via API calls to the IoT resource, wherein an API for the API call is included in the IoT resource description.

19. The system of claim 16, wherein the data about the IoT resource comprises privacy requests supported by the IoT resource.

20. The system of claim 16, wherein the data about the IoT resource comprises identification of documentation that demonstrates qualification to submit the privacy request.

21. The system of claim 16, wherein the data about the IoT resource comprises default privacy practices of the IoT resource.

22. The system of claim 16, wherein the PPA receives the data about the IoT resource from an IoT resource description for the IoT resource hosted by an IoT privacy infrastructure, and wherein the IoT resource description: identifies one or more authorized agents that are trusted sources for privacy requests for the IoT resource; and APIs to communicate with the one or more authorized agents.

23. The system of claim 16, wherein the PPA is further configured to determine the privacy request for the user using a privacy preference model for the user and using the data about the IoT

resource.

24. The system of claim 16, wherein the PPA is further configured to, in response to a request, communicate to the authorized agent electronic documentation that demonstrates that the user qualifies to submit the privacy request.

25. The system of claim 16, further comprising an IoT resource registry for advertising the IoT resource, and wherein the PPA is for discovering the IoT resource registry based on a location of the computing device of the user.

26. A method for configuring a privacy request for an Internet of Things (IoT) resource that is for remotely sensing data about a user, the method comprising: receiving, by a personal privacy app (“PPA”) running on a computer device of the user, data about the IoT resource, wherein the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; communicating a privacy request for the user with respect to the IoT resource, wherein the privacy request is one of the available user-specific privacy requests and wherein the privacy request communicated by the PPA is based on the data received about the IoT resource; and in response to a query related to the privacy request, causing electronic documentation to be transmitted that demonstrates that the user qualifies to submit the privacy request, such that the privacy request is applied to data collected about the user by the IoT resource.

27. A method for configuring a privacy request for an Internet of Things (IoT) resource that is for remotely sensing data about a user, the method comprising: receiving, by a personal privacy app (“PPA”) running on a computer device of the user, data about the IoT resource, wherein the data about the IoT resource comprises available user-specific privacy requests related to data practices of the IoT resource; communicating, by the PPA via an authorized agent, a privacy request for the user with respect to the IoT resource, wherein the privacy request is based on the data received about the IoT resource, and wherein the authorized agent is a trusted source of privacy requests for the IoT resource; and communicating, by the authorized agent, the privacy request for the user for implementation with respect to data collected about the user by the IoT resource.
