

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication	20250265364
Kind Code	A1
Publication Date	August 21, 2025
Inventor(s)	Jassal; Amrit et al.

Hybrid Approach To Data Governance

Abstract

A cloud-based data governance system includes a processing unit, a network adapter, and memory for storing data and code. The network adapter establishes a connection with a remote data storage system associated with a remote file system over a wide-area network (WAN). The code includes and event collection interface, a data governance service, and an enforcement service. The event collection interface is configured to capture an event from the remote data storage system. The event is indicative of a file system operation executed on a data object of the remote file system. The data governance service is configured to receive the event from the event collection interface and to process the event to determine whether the file system operation conflicts with a governance policy of the data governance system. The enforcement service executes a set of remediation actions, if the file system operation does conflict with the governance policy.

Inventors: Jassal; Amrit (Mountain View, CA), Sharma; Shishir (Mountain View, CA), Puttergill; Sean H. (Sunnyvale, CA), Sundararaj; Ramakrishnan (Sunnyvale, CA)

Applicant: Egnyte, Inc. (Mountain View, CA)

Family ID: 1000008578174

Assignee: Egnyte, Inc. (Mountain View, CA)

Appl. No.: 19/027189

Filed: January 17, 2025

Related U.S. Application Data

parent US continuation 18802630 20240813 PENDING child US 19027189

parent US continuation 17960546 20221005 parent-grant-document US 12093417 child US 18802630

parent US continuation 15487947 20170414 parent-grant-document US 11494503 child US

Publication Classification

Int. Cl.: **G06F21/62** (20130101); **G06F16/11** (20190101); **G06F16/182** (20190101); **G06F16/23** (20190101)

U.S. Cl.:

CPC **G06F21/6218** (20130101); **G06F16/122** (20190101); **G06F16/128** (20190101); **G06F16/183** (20190101); **G06F16/2365** (20190101);

Background/Summary

RELATED APPLICATIONS [0001] This application is a continuation of co-pending U.S. patent application Ser. No. 18/802,630, filed on Aug. 13, 2024 by the same inventors, which is a continuation U.S. patent application Ser. No. 17/960,546, filed on Oct. 5, 2022 by the same inventors, which is a continuation of U.S. patent application Ser. No. 15/487,947, filed on Apr. 14, 2017 by the same inventors, which claims the benefit of priority to U.S. Provisional Patent Application Ser. No. 62/322,722, filed on Apr. 14, 2016 by the same inventors. All prior applications are incorporated herein by reference in their respective entireties.

BACKGROUND

Field of the Invention

[0002] This invention relates generally to cloud computing systems, and more particularly to data governance for cloud computing systems.

Description of the Background Art

[0003] Data governance applications are known. Data governance applications allow a user to track data access and modification events across a variety of sources. Some applications allow a user to define and enforce policies to ensure adequate data security. Typically, a data governance application is enterprise-based software that is installed and operated on-premises for tracking access and modification records for each data object on one or more local servers.

[0004] As cloud-based data storage and computing have become more popular, the need for data governance for cloud-based data has also increased. One solution for tracking events on a cloud-based server is the utilization of on-premises data governance at each local site associated with a particular cloud server. This solution is disadvantageous, because utilizing data governance deployed on a local system to govern data usage of a remote system creates security concerns. In addition, because these applications are localized, the amount of data that can be processed is limited by the physical properties of the local network on which they are deployed.

[0005] Another solution is the utilization of data governance applications on the cloud itself. This solution addresses security concerns, as the data governance application is no longer deployed remotely from the data sources. However, this solution does not provide for data governance of data sources that are deployed locally. It is necessary for separate data governance applications to be deployed on the cloud and on-premises, increasing cost and introducing the need for interfacing between the separate data governance applications. What is needed, therefore, is a single data governance application for monitoring local and cloud-based data sources.

SUMMARY

[0006] The present invention overcomes the problems associated with the prior art by providing a

cloud-based means for delivering data governance services to a multitude of data sources located on one or more client premises and/or other cloud-based data sources remote with respect to the cloud-based data governance system. The invention facilitates the provision of data governance services via cloud-based software as a service (SaaS).

[0007] Methods that can be implemented, for example, in a cloud-based data governance system are disclosed. An example method for providing data governance of a remote data storage system associated with a remote file system includes establishing a connection with the remote data storage system over a wide area network (WAN) and capturing an event associated with the remote file system. The event is indicative of at least one file system operation executed on a data object of the remote data storage system. The example method further includes processing the event to determine whether the event conflicts with a governance policy of the data governance system and, if the event does conflict with the governance policy, executing a set of remediation actions.

[0008] In a particular example method, the step of capturing an event associated with the remote file system includes deploying an event collection service to the remote data storage system. The event collection service is operative to detect file system operations executed on data objects of the remote data storage system, to generate events indicative of the file system operations, and to push the events to the data governance system. The method further includes receiving the events from the remote data storage system via the event collection service.

[0009] Another particular example method additionally includes receiving a metadata snapshot from the remote data storage system. The metadata snapshot is indicative of the remote file system, and the example method also includes generating a derivative data set indicative of the remote file system based on the metadata snapshot. The step of capturing an event associated with the remote file system includes capturing metadata associated with one or both of at least one file system operation and a data object of the file system. Optionally, the step of capturing metadata includes capturing metadata indicative of a particular user executing the at least one file system operation, and the step of executing a set of remediation actions includes altering permissions associated with the particular user.

[0010] In a particular method, the step of processing the event includes creating and/or updating a derivative data set based on the event. The derivative data set is derived from the data of the remote data storage system associated with the remote file system. The method further includes performing data analytics on the derivative data set after the derivative data set has been updated. Optionally, the step of processing the event includes performing data analytics on the event itself.

[0011] In one example method, the step of executing a set of remediation actions includes pushing a control message to the remote data storage system. The control message indicates a set of file system operations to be executed on objects of the remote file system by the remote data storage system.

[0012] Example methods operate on a continuous basis, collecting additional events. Each event of the additional events is indicative of at least one additional file system operation executed on a data object of the remote file system stored on the remote data storage system. The method additionally includes storing the event and the additional events in an event database and providing a client associated with the remote file storage system access to the event database.

[0013] Example methods also facilitate data governance of data sources stored on third-party storage systems, remote with respect to both a client's site and the data governance site. In one method, the step of establishing a connection with the remote data storage system includes establishing a connection with a third party cloud service provider separate from the cloud-based data governance system.

[0014] An example cloud-based data governance system is also disclosed. The cloud-based data governance system includes a processing unit, a network adapter, and memory for storing data and code. The processing unit is configured to execute the code to impart functionality to the system. The network adapter is electrically coupled to establish a connection with a remote data storage

system associated with a remote file system over a wide-area network (WAN).

[0015] The code includes an event collection interface, a data governance service, and an enforcement service. The event collection interface is configured to capture an event from the remote data storage system. The event is indicative of at least one file system operation executed on a data object of the remote file system stored on the remote data storage system. The data governance service is configured to receive the event from the event collection interface and to process the event to determine whether the at least one file system operation conflicts with a governance policy of the data governance system. The enforcement service is configured to execute a set of remediation actions, if the at least one file system operation does conflict with the governance policy.

[0016] In a particular example embodiment, the event collection interface is configured to deploy an event collection service to the remote file storage system. The event collection service is operative to detect file system operations executed on data objects of the remote file system stored on the remote data storage system and to generate events indicative of the file system operations. The event collection service then pushes the events to the data governance system. The event collection interface is configured to receive the events from the remote data storage system via the event collection service. Optionally, the event collection interface can periodically poll the event collection service for the events.

[0017] In an example embodiment, the event collection interface is further configured to receive a metadata snapshot of the remote data storage system. The metadata snapshot is indicative of the remote file system, and the data governance service is further configured to generate a derivative data set based on the metadata snapshot. The event collection interface is configured to capture metadata associated with file system operation(s) and/or data object(s) associated with captured event(s). The data governance service is additionally configured to update the derivative data set based on the captured event(s) and to perform data analytics on the updated derivative data set.

[0018] The enforcement module is additionally configured to push one or more control messages to the remote data storage system, if a data governance policy is violated. The control message(s) indicate a set of file system operations to be executed on objects of the file system on the remote data storage system. As a non-limiting example, the event collection interface can be configured to capture metadata indicative of a particular user executing the file system operation, and the set of remediation actions can include altering permissions associated with the particular user.

[0019] The system can further include an event database operative to store records of the captured events. The event collection interface is configured to collect additional events and store records of the additional events in the database. Each event of the additional events is indicative of at least one additional file system operation executed on a data object of the remote file system stored on the remote data storage system. Events can also be generated by monitoring update events associated with other data source types (e.g., Egnyte Connect, SharePoint, Windows Server, etc.) associated with other subsystems (e.g., file systems, links, permissions, etc.) that are present. Optionally, a client interface is configured to provide a client associated with the remote file system access to the event database. As another option, the data governance service is additionally configured to perform batch data analysis functions on a subset of the records of the database. As another option, the data governance service can be additionally configured to perform data analytics on the individual events.

[0020] In a particular embodiment, the remote computer system is a third party cloud service provider.

[0021] Methods that can be implemented, for example, in a local data storage system are also disclosed. An example method for utilizing cloud-based data governance services includes capturing an event indicative of a file system operation performed on a data object stored in the local data storage system. The method additionally includes establishing a connection with a remote cloud-based data governance system over a wide-area network (WAN) and providing the

event to the data governance system. Provision of the event facilitates a determination of whether the event conflicts with a data governance policy stored on the data governance system. The method additionally includes executing a set of remediation actions on the local data storage system responsive to one or more communications from the data governance system, if the data governance system determines that the event conflicts with a data governance policy stored on the data governance system.

[0022] In a particular example method, the step of capturing an event includes deploying a plurality of data monitors. Each of the plurality of data monitors is associated with one of a plurality of different data source types, and each of the data monitors is operative to detect file system operations executed on data objects of the associated data source type. The data monitors then generate events indicative of the file system operations, and push the events to the data governance system.

[0023] Example methods of capturing/generating events are disclosed. In one example method the step of generating events includes scanning at least one of the data source types at different times. In another example method, the step of generating an event includes registering for callbacks from an application associated with at least one of the data source types. In yet another example method, the step of generating an event includes intercepting and filtering events from at least one of the data source types. Optionally, the steps of intercepting and filtering events from the at least one data source type includes installing an agent on-site with the local data storage system, the agent being configured to intercept and filter the events. The step of generating an event can include capturing metadata associated with the file system operation and/or the data object. In a particular example method, the step of capturing metadata includes capturing metadata identifying a particular user performing the file system operation on the data object. These example methods of capturing/generating events, as well as others, can be used individually or in any combination with one another, as the needs of a particular application might dictate.

[0024] In an example method, the step of providing the event to the data governance system includes providing metadata of a file system associated with the local data storage system to facilitate the creation or updating of a derivative data set by the data governance system. The step of providing the event to the data governance system can also include providing at least a portion of the data object associated with the event to facilitate the creation or updating of a derivative data set by the data governance system.

[0025] In an example method, the step of executing a set of remediation actions includes receiving one or more control messages indicating the set of remediation actions to be executed on the local data storage system. In a particular example method the step of executing a set of remediation actions on the local data storage system includes altering permissions associated with a particular user identified by an event.

[0026] A local data storage system is also disclosed. The local data storage system includes a processing unit configured to execute code, a network adapter electrically coupled to establish a connection with a remote cloud-based data governance system over a wide-area network (WAN), and memory. The memory stores data and the code. The data and the code include an event collection service, a data governance interface, and an enforcement module. The event collection service is configured to capture an event, which is indicative of a file system operation performed on a data object of the local data storage system. The data governance interface is configured to provide information associated with the event to the data governance system. The enforcement module is responsive to communications from the data governance system and is operative to execute a set of remediation actions on the local data storage system, if the data governance system determines that the event creates a conflict with a data governance policy stored on the data governance system.

[0027] In an example system, the event collection service includes a plurality of data monitors. Each of the plurality of data monitors is associated with one of a plurality of different data source

types. Each data monitor is also operative to detect file system operations executed on an associated data source of the associated type, to generate events indicative of the file system operations, and push the events to the data governance system.

[0028] Various example data monitor functions are disclosed. As one example, at least one of the data monitors is configured to scan the associated data source at different times in order to detect the file system operations. As another example, at least one of the data monitors is configured to register for callbacks from an application associated with the associated data source type in order to detect the file system operations. As yet another example, at least one of the data monitors is configured to intercept and filter events related to the associated data source in order to detect the file system operations. As yet another example, at least one of the data monitors is configured to install an agent on the particular data source, the agent being configured to intercept and filter the events. The data monitors can be additionally configured to capture metadata associated with one or both of the file system operation and the data object. In an even more detailed example, the metadata is indicative of a particular user performing the file system operation on the data object. Any or all of the disclosed data monitor functions can be used in any combination with each other or with other data monitor functions depending on the needs of a particular system.

[0029] In an example system, the information associated with the event, which is provided to the data governance system by the data governance interface, includes metadata of a file system associated with the data object to facilitate the creation or updating of a derivative data set by the data governance system. Optionally, the information associated with the event includes at least a portion of the data object to facilitate the creation or updating of the derivative data set by the data governance system.

[0030] An example system includes remediation capabilities. The data governance interface is configured to receive one or more control messages indicating a set of remediation actions to be executed on the local file storage system. The enforcement module is configured to execute the set of remediation actions responsive to the one or more control messages. In a particular example embodiment, the remediation actions include altering permissions associated with a particular user identified by an event.

[0031] An example event collection system is also disclosed. The example event collection system is deployable on a file storage system and includes a processor, a network adapter, and memory. The processor is configured to execute code. The network adapter is electrically coupled to establish a connection to a data governance service over a wide-area network (WAN). The data governance service is located remotely from the event collection system. The memory provides storage for data and code. The data and code include a source connector routine and a data governance interface. The source connector routine is configured to monitor a corresponding particular data source on the file storage system and to generate an event responsive to a file system operation being executed on a data object associated with the particular data source. The event is indicative of the file system operation. The data governance interface is configured to push the event to the data governance service.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0032] The present invention is described with reference to the following drawings, wherein like reference numbers denote substantially similar elements:

[0033] FIG. 1 is a diagram showing an example cloud-based data governance system;

[0034] FIG. 2 is a relational diagram showing data transfer between a data governance system of FIG. 1 and a local/remote data source;

[0035] FIG. 3A shows an example of event communication between a client site and a remote data

governance system of FIG. 1;

[0036] FIG. 3B shows another example of event communication between a client site and a remote data governance system of FIG. 1;

[0037] FIG. 3C shows yet another example of event communication between a client site and a remote data governance system of FIG. 1;

[0038] FIG. 4 is a diagram showing the client site of FIG. 1 in greater detail;

[0039] FIG. 5 is a block diagram showing the connector hub host device of FIG. 4 in greater detail;

[0040] FIG. 6 is a block diagram showing various aspects of the connector hub of FIG. 5;

[0041] FIG. 7 is a block diagram showing the connector hub of FIG. 5 in even greater detail, including specific examples of the connectors, applications, and platform services of FIG. 6;

[0042] FIG. 8 is a block diagram showing the data governance system of FIG. 1 in greater detail;

[0043] FIG. 9 is a block diagram showing aspects of the data governance services of FIG. 8 in greater detail;

[0044] FIG. 10A is a data flow diagram showing a method for initially configuring the data governance system and the local file storage system of FIG. 1 for data governance services;

[0045] FIG. 10B is a data flow diagram showing a method for initially configuring the data governance system and the storage server of FIG. 1 for data governance services;

[0046] FIG. 10C is data flow diagram showing a method for updating derivative data sets stored on the data governance system of FIG. 1;

[0047] FIG. 11 is a data flow diagram showing a portion of the method of FIG. 10C in greater detail;

[0048] FIG. 12 is a data flow diagram showing the connector hub of FIG. 5 retrieving events, metadata, and content from data sources and forwarding the events, metadata, and content to the data governance system of FIG. 1;

[0049] FIG. 13 is a stacked software diagram of the connector hub of FIG. 5, including example ones of the connectors, applications, and platform services of FIG. 6;

[0050] FIG. 14 is a process flow diagram showing an example method for retrieving a snapshot of a source of FIG. 12 via an agent installed on the source;

[0051] FIG. 15A is a process flow diagram showing an example method for processing a folder snapshot in the data governance server of FIG. 8;

[0052] FIG. 15B is a process flow diagram showing an example method for processing a file snapshot in the data governance server of FIG. 8;

[0053] FIG. 16 is a process flow diagram showing a particular example method for providing a directory service snapshot to the data governance server of FIG. 8;

[0054] FIG. 17 is a process flow diagram showing an example method for scanning a source of FIG. 12 to generate events;

[0055] FIG. 18 is a process flow diagram showing an example method for intercepting and filtering events via an agent on a source of FIG. 12;

[0056] FIG. 19 is a process flow diagram showing a particular example method for intercepting and filtering events via the agent of FIG. 18;

[0057] FIG. 20 is a process flow diagram showing an example method for receiving events from the source of FIG. 12 by registering for callbacks;

[0058] FIG. 21 is a flow chart summarizing an example method of performing data governance from a site remote from a data source;

[0059] FIG. 22 is a flow chart summarizing an example method of performing the fourth step of the method of FIG. 21;

[0060] FIG. 23 is a flow chart summarizing a method of utilizing data governance services from a remote site; and

[0061] FIG. 24 is a flow chart summarizing an example method of performing the fourth and fifth steps of the method of FIG. 23.

DETAILED DESCRIPTION

[0062] The present invention overcomes the problems associated with the prior art, by providing a cloud-based data governance system hosted on a remote computer system with respect to the data source being governed. The data governance system provides data governance services, including data analytics, for data sources hosted on a local file system (e.g., a client site) or a remote, cloud-based storage system. The data governance system deploys event collection software to the local file system, to collect events generated in response to access, modification, and/or other operations executed on data sources of the local file system, as well as metadata and content. The data governance system also utilizes publicly available application programming interfaces (APIs) to collect access, modification, and/or other events from the cloud-based storage system. The data governance system analyzes the events in order to detect risks, threats, suspicious behavior, or any condition that violates an existing data governance policy, and utilizes the event collection software and APIs to send remediation operations to the local file system and the cloud-based storage system, respectively, in order to provide data security.

[0063] In the following description, numerous specific details are set forth (e.g., data types, event types, protocols, etc.) in order to provide a thorough understanding of the invention. Those skilled in the art will recognize, however, that the invention may be practiced apart from these specific details. In other instances, details of well-known cloud computing practices (e.g., virtualization, load balancing, etc.) and components have been omitted, so as not to unnecessarily obscure the present invention.

[0064] FIG. 1 is a diagram showing an example cloud computing system **100** that includes a cloud-based, hybrid data governance system **102**, a local file storage system **104**, and a cloud-based storage server **106**, all interconnected via the Internet **108**. Data governance system **102** provides visibility to data access and modification events occurring on local file storage system **104** and data storage server **106**. Data governance system **102** can be accessed by remote users logged onto remote clients **110** via Internet connections **112** or alternative connections (e.g. dial-up connections) **114**. Remote clients **110** are machines (e.g. laptops, smart phones, etc.) with sufficient credentials to view data indicative of file system operations that are executed on data objects stored on local file storage system **104** and storage server **106**. Remote users can also utilize remote clients **110** to view and/or change governance policies stored on data governance system **102**.

[0065] Local file storage system **104** can be hosted, for example, on a network-attached storage (NAS) device (FIG. 2) on a local network **116** located at a client site **118(1)**. Additional client sites **118(2-C)** host additional local file storage systems. Each of client sites **118(1-C)** can be associated with the same or different cloud clients, as data governance system **102** is capable of providing governance services to any number of clients at any number of different locations. Local users can utilize local clients **120** to access and/or modify data objects stored on local file storage system **104** and also, with proper credentials, view and/or change governance policies stored on data governance system **102**. In the example embodiment, at least a portion of local file storage system **104** is bi-directionally synchronized with storage server **106**. In alternate embodiments, local file storage system **104** and storage server **106** can operate completely independently of one another. Storage server **106** is a cloud-based application for storing and accessing remote data objects. Remote clients **110** can access storage server **106** via Internet connections **112** or alternative connections **122**, in order to upload, download, view, or update data objects stored thereon. Optionally, local clients **120** can also access storage server **106** via local network **116** and Internet **108**.

[0066] FIG. 2 is a relational diagram showing data transfer between data governance system **102** and a local/remote data source **202**. Local/remote data source **202** is defined with respect to local file storage system **104** and can be a data source stored thereon or a data source located on a remote service, such as storage server **106**. In either case, data source **202** is located remotely from data governance system **102** and communicates bi-directionally with data governance system **102**. Data

source **202** also sends events, metadata, and content to data governance system **102**. Events include notifications that a data object on data source **202** was accessed or edited in some way. Metadata includes data representative of the file system, the file system directory, and permissions associated with file system objects on local file storage system **104** and/or cloud-based storage **106**. Content includes data objects themselves, for example a WORD document, EXCEL file, etc. Data governance system **102** requests, receives, and processes the events, metadata, and content in order to provide data governance services for data source **202**. Additionally, data governance system **102** sends control messages, including, but not limited to, commands to execute file system operations, to data source **202**.

[0067] FIGS. **3A-3C** are relational diagrams showing data transfer (event communication) between data governance system **102** and various data sources, each shown in a separate example system. FIG. **3A** shows an example data source **302**, hosted on client site **118(1)**, in communication with a hub **304**. Hub **304** receives events, metadata, and content directly from data source **302**. Hub **304** maintains an Internet connection with a hub interface **306** on data governance system **102** and sends the events, metadata, and content from data source **302** to hub interface **306** via the connection. Hub **304** and hub interface **306** each include specific networking protocols for communicating with one another over the Internet. Hub interface **306** forwards the data (e.g., events, metadata, and/or content) received from hub **304** onto a governance application **308**, which performs data analytics and provides other data governance services based on the received data.

[0068] FIG. **3B** shows an example data source **310** hosted on client site **118(1)**. Data source **310** is substantially similar to data source **302**, except data source **310** and hub **304** cannot directly communicate with one another, at least for some event types. Therefore, a source agent **312** is also hosted on client site **118(1)**. Source agent **312** is a software module that provides an interface between data source **310** and hub **304**. It should be noted that although source agent **312** is shown separately from data source **310**, in reality source agent **312** could be installed directly onto data source **310**. Hub **304**, hub interface **306** and governance application **308** function as described with respect to FIG. **3A**.

[0069] FIG. **3C** shows an example data source **314** hosted on storage server **106**. Data source **314** utilizes one or more application programming interfaces (APIs) **316** to facilitate communication with its clients via the Internet **108**. Cloud connectors **318** utilizes APIs **316** to facilitate communication between data governance system **102** and storage server **106**. APIs **316** include publicly available protocols for communicating with remote services over the Internet. Cloud connectors **318** utilize APIs **316** to retrieve events, metadata, and content from storage server **106** for data governance system **102**. Cloud connectors **318** additionally include software for generating events based on the data retrieved from storage server **106**. Cloud connectors **318** then forward events, metadata, and content received via APIs **316** onto governance application **308**.

[0070] FIG. **4** is a block diagram showing communication between various components of cloud computing system **100**, including client site **118(1)**, which is shown in greater detail. Client site **118(1)** includes a network-attached storage (NAS) device **402**, a WAN adapter **404**, a connector hub host device **406**, and local clients **120**, all interconnected via local network **116**. NAS device **402** is a storage device connected to local network **116** and accessible by other components connected to local network **116**. NAS device **402** hosts data source(s) **408**, and a directory service **410** runs on a separate, dedicated server. Data sources **408** include file system objects (e.g. files, metadata, applications, etc.) constituting a local file system that can be accessed by local clients **120** for viewing, editing, utilization, etc. Directory service **410** includes user permissions and lookup tables to allow local clients **120** with sufficient credentials to locate and access available data objects included in data sources **408**. WAN adapter **404** is a network device that provides a connection to a wide-area network, which, in this example, is the Internet **108** (omitted from FIG. **4** for clarity). Components connected to local network **116** can access data governance system **102** and storage server **106** via an Internet connection **412** provided by WAN adapter **404**. Local clients

120 can utilize Internet connection **412** to upload and/or download data objects from storage server **106**.

[0071] Connector hub host device **406** is a device that hosts a software-based connector hub (FIG. 5). In the example embodiment, connector hub host device **406** is a server hosting virtualization software for running virtual machines to host various components of the connector hub. The connector hub monitors data sources **408** including directory service **410** to detect access and modification to data objects by local clients **120**. The connector hub generates events based on the access and modifications detected, and connector hub host device **406** sends the events to data governance system **102** via WAN adapter **404** and Internet connection **412**. Additionally data governance system **102** receives and/or pulls access and modification events from storage server **106**, resulting from access and modification by local clients **120** or remote clients **110** (FIG. 1).

[0072] FIG. 5 is a block diagram showing connector hub host device **406**, including a connector hub **502**, in greater detail. Connector hub host device **406** is a server for hosting virtualized machines and includes non-volatile memory **504**, one or more processing units **506**, working memory **508**, a local network adapter **510**, one or more user interface devices **512**, connector hub **502**, and a governance enforcement module **514**, all interconnected via a system bus **516**. Non-volatile memory **504** is a data storage device that stores data objects, such as files and software, to be accessed by other elements of connector hub host device **406** and local network **116**. Non-volatile memory **504** can include several different storage devices and types, including hard disk drives, solid state drives, read-only memory (ROM), etc. distributed across local network **116**. Processing unit(s) **506** transfer code from non-volatile memory **504** into working memory **508** and execute the code to impart functionality to various components of connector hub host device **406**. For example, working memory **508** stores code, such as software modules, that when executed provides the described functionality of connector hub **502**. Local network adapter **510** provides a network connection between connector hub host device **406** and local network **116** and, therefore, WAN adapter **404** and the Internet **108** (FIGS. 1 and 4). User interface device(s) **512** (e.g. keyboards, mice, etc.) enable local IT personnel to access connector hub host device **406**, e.g., for firmware upgrades, software upgrades, etc.

[0073] Connector hub **502** is a framework of virtualized nodes for generating data access and modification events and sending the events to data governance system **102**. Connector hub **502** monitors data sources **408** and directory service **410** (FIG. 4) and generates events responsive to and indicative of access, modifications, and other operations executed on directory service **410** and the local file system stored on data sources **408**. These events are sent to data governance system **102** via local network adapter **510** and WAN adapter **404**, in order to be processed and stored for data governance analytics and visibility purposes. Governance enforcement module **514** is a virtualized software module that receives control messages from data governance system **102**. Governance enforcement module **514** can access and alter data sources **408** and directory service **410**, responsive to receiving control messages from data governance system **102**, in order to quarantine suspicious files or alter permissions for a user engaging in suspicious activities, by way of non-limiting example. Together, connector hub **502** and governance enforcement module **514** constitute a framework for collecting events from and enforcing data governance policies on the local file system.

[0074] FIG. 6 is a block diagram showing various components of connector hub **502**. Connector hub **502** includes a set of connectors **602(1-N)**, a set of applications **604(1-M)**, and a set of platform services **606(1-P)**. Connectors **602**, applications **604**, and platform services **606** are all software modules running in working memory **508** on a hub system layer. Connectors **602** are included in a connector layer **610**, and each of connectors **602** monitor a specific data source, where a data source includes a subset of data sources **408** that corresponds to a particular data type and/or program. For example, a source can be word processing software, such as Microsoft Word, and associated data, such as documents, settings, system files, etc. Connectors **602** monitor the

corresponding sources and generate events whenever the respective sources are accessed or modified. These events provide data governance system **102** with information about the file system object, the file system operation executed thereon, and the user that executed the operation. Additionally, connectors **602** can send copies of content (e.g., part or all of a file system object) for further analysis, when such further analysis is suggested by a data governance policy.

[0075] Applications **604** are included in an application layer **612**, and each of applications **604** provide a separate service for analyzing events created by connectors **602**. Applications **604** provide services that can be utilized, for example, to determine that a file system object contains malware or that a user is downloading sensitive material, based on an event, multiple events, or the file system objects themselves. More specific examples of applications **604(1-M)** will be provided herein with reference to subsequent drawings.

[0076] Platform services **606** are included in a platform layer **614**, and each of platform services **606** provides connectors **602** and applications **604** with underlying, support functionality, such as communication with hardware devices. For example, platform services **606** allow the other components of connector hub **502** to communicate with one another and with devices hosted on connector hub host device **406**, such as local network adapter **510**.

[0077] Connector layer **610**, application layer **612**, and platform layer **614** constitute a layered, software framework that is more fully described with reference to FIG. **13** below.

[0078] FIG. **7** is a block diagram showing connector hub **502** in even greater detail, including specific examples of connectors **602**, applications **604**, and platform services **606**. Connectors **602** include directory service connector(s) **702**, content management connector(s) **704**, messaging connector(s) **706**, file sync & share connector(s) **708**, email connector(s) **710**, and file server connector(s) **712**. Directory service connector(s) **702** monitor directory service **410** for changes made thereto. For example, if one of local clients **120** is disconnected from local network **116**, a directory service connector **702** will generate an event indicative of the corresponding change to the data contained in directory service **410**. Content management connector(s) **704** monitor content management software, such as Microsoft Sharepoint, EMC Documentum, etc., and associated data. Messaging connector(s) **706** monitor messaging software, such as Twitter, Yahoo Instant Messenger, Slack, etc., and associated data. File sync and share connector(s) **708** monitor file sharing software, such as Egnyte Connect, Box, Dropbox, etc., and associated data. Email connector(s) **710** monitor email software, such as Microsoft Exchange, Google Gmail, etc., and associated data. File server connector(s) **712** monitor file server software, such as Netapp, EMC, Dell, etc., and associated data.

[0079] Applications **604** include content extraction service(s) **714**, content detection service(s) **716**, and pattern detection service(s) **718**. Content extraction service(s) **714** analyze data and content to extract data into a text file or a PDF. For example, a content extraction service **714** converts software code into a parse-able and query-able text file, such as a notepad document. Content detection service(s) **716** analyze the extracted content to determine parts, sub-parts, and MIME types of the content. For example, a content detection service **716** might analyze the extracted notepad document to determine that the underlying content is an HTML document with Java Plugins. Pattern detection service(s) **718** analyze the extracted notepad document to detect pre-defined patterns indicative of relevant issues (e.g., security, privacy, and so on) from a data governance perspective. For example, a predefined pattern might include numbers grouped in a ###-##-#### pattern, which could be indicative of an employee's social security number.

[0080] Platform services **606** include source-type connectors **720**, a lock manager **722**, a distributed configuration **724**, a message queue **726**, a distributed cache **728**, a transport service **730**, and a discovery service **732**. Source-type connectors **720** are generic connectors by source type for retrieving additional metadata from sources such as file systems. Source-type connectors **720** can, for example, retrieve file attributes related to specific events. Lock manager **722** is a service that ensures sequential access to critical resources or locks critical resources for performing operations,

such as configuration changes across nodes. Distributed configuration **724** is a resilient service that shares configuration data, such as source credentials, securely across nodes. Message queue **726** is a resilient queue used to push data and events between the various nodes. For example, directory service connector **702** can use message queue **726** to provide data or an event to content extraction service **716** for analysis. Distributed cache **728** is distributed memory used for storing frequently looked up data, such as mapping of security identifiers to user information. Transport service **730** manages a pool of always-on, bi-directional connections, such as WebSocket connections, to data governance system **102**. Discovery service **732** allows each of connectors **602**, applications **604**, and platform services **606** to dynamically discover each other across the various nodes. It should be noted that, in alternate embodiments, any and all of connectors **602**, applications **604**, and platform services **606** can be hosted by data governance system **102**, instead of connector hub **502**, or be distributed between the two. Alternatively, any of connectors **602**, applications **604**, and platform services **606** can be replicated in data governance system **102**, as is shown in FIG. **9** below.

[0081] FIG. **8** is a block diagram showing data governance system **102** in greater detail. Data governance system **102** is a cloud-based computer system including multi-tenant data storage devices **802**, a WAN adapter **804**, and data governance servers **806(1-S)**, all interconnected via a local network **808**. Storage devices **802** are network attached storage devices for storing data associated with multiple different cloud clients. Storage devices **802** can provide the non-volatile data storage utilized by every other component of data governance system **102**. WAN adapter **804** is a network adapter for establishing a connection to the Internet **108**. Elements of data governance system **102** utilize WAN adapter **804** to communicate with remote systems, such as local file storage system **104** (e.g., connector hub **502**) and storage server **106**.

[0082] Data governance servers **806** provide data governance services for local file storage systems and cloud-based storage servers associated with various cloud clients. In the example embodiment, data governance server **806(1)** provides data governance services for local file storage system **104** and storage server **106**. Data governance server **806(1)** includes one or more processing units **810(1)**, working memory **812(1)**, a local network adapter **814(1)**, and a data governance services module **816(1)**, all interconnected via an internal bus **818(1)**. Processing unit(s) **810(1)** execute code transferred into working memory **812(1)** from, for example, storage devices **802**, to impart functionality to various components of data governance server **806(1)**. Working memory **812(1)** can also cache frequently used code, such as network locations of storage devices **802**, to be quickly accessed by the various components of data governance server **806(1)**. Local network adapter **814(1)** provides a network connection between data governance server **806(1)** and local network **808** and, therefore, WAN adapter **804**, which provides a connection to the Internet **108**. Data governance services **816(1)** are various software services, running within working memory **812(1)**, for collecting and analyzing events that are received from connector hub **502**. Data governance services **816(1)** perform data analytics on events and file system metadata received from connector hub **502**. Although only data governance server **806(1)** is shown in detail, it should be understood that data governance server **806(1)** is substantially similar to data governance servers **806(2-S)**, except that any of data governance servers **806** can correspond to different cloud clients and, therefore, can be configured differently to utilize different data, connectors, applications, network connections, etc.

[0083] FIG. **9** is a block diagram showing components of data governance services **816(1)** in greater detail. Data governance services **816(1)** perform several functions, including receiving and processing events, metadata, and content from local file storage system **104** and storage server **106**, generating and updating a derivative data set, performing data analytics, providing visibility to data access and events, and providing control messages to local file storage system **104** and storage server **106**.

[0084] Data governance services **816(1)** receive events, metadata, and/or content from local file storage system **104** and/or storage server **106** via the Internet. A connector hub interface **902**

receives messages containing events, metadata, or content from connector hub **502** and removes any protocol headers (e.g., WebSocket headers) before saving the messages in an incoming queue **904**. Similarly, one or more cloud connectors **906**, which are configured to communicate with cloud-based data sources by utilizing publicly available APIs, retrieve information from storage server **106**, generates events based on the information, and stores the events in incoming queue **904**. A message processor **908** reads the events from incoming queue **904** and determines whether the message constitutes an event, metadata, or content. Message processor **908** saves events in an event store **910**, saves metadata in metadata store **912**, and saves content in an object store **914**. Whenever message processor **908** processes a message, it also notifies a governance services manager **916** that an event, metadata, or content has been received and processed. In response, governance services manager **916** uses events stored in events store **910** or metadata stored in metadata store **912** to generate (for the first time) or update a derivative data set, which includes metadata indicative of local file storage system **104** (or storage server **106**).

[0085] Governance services manager **916** utilizes a set of governance services **918** to perform data processing and analytics on incoming events, metadata, and content, as well as the derivative data set stored in event store **910**, metadata store **912**, and/or object store **914**. Governance services **918** include a snapshot service **920**, a file system service **922**, a directory service **924**, a permissions service **926**, an events service **928**, a content extraction service **930**, a content detection service **932**, and a pattern detection service **934**. Snapshot service **920** coordinates file system service **922**, directory service **924**, and permissions service **926**, in order to capture metadata and edit the derivative data set on metadata store **912**. Snapshot service **920** controls when (e.g. every 10 minutes) and how to capture metadata from local file storage system **102**. File system service **922** selects file system data from the received metadata and generates/alters a cloned file system tree indicative of the file system on local file storage system **104** as part of the derivative data set. Directory service **924** selects directory data from the received metadata and generates/alters a cloned directory tree indicative of the directory tree on local file storage system **104** as part of the derivative data set. Permissions service **926** selects permissions data from the received metadata and generates/alters a cloned permissions tree indicative of the permissions tree on local file storage system **104** as part of the derivative data set. Events service **928** analyzes events stored in event store **910** in order to determine whether or not to modify the derivative data set stored in metadata store **912** or if additional metadata or content from local file storage system **104** is needed. If additional metadata or content is required, governance services manager **916** can request the necessary data from local file storage system **104**, as will be described below. Together, services **920**, **922**, **924**, **926**, and **928** provide the functionality required to generate and/or update the entire derivative data set stored in metadata store **912**. Content extraction service **930**, content detection service **932**, and pattern detection service **934** are substantially similar to content extraction service **714**, content detection service **716**, and pattern detection service **718**, of connector hub **502**.

[0086] Utilizing governance services **918**, governance services manager **916** analyzes events on event store **910**, the derivative data set on metadata store **912**, and objects on object store **914**, in view of a set of governance policies **936**. Governance policies **936** include a vast set of predefined criteria including, but not limited to, security criteria, privacy criteria, definitions of suspicious and/or threatening activity and data, including patterns indicative of malicious code and system attackers, access criteria, and so on. Governance policies **936** also include remediation definitions, which provide governance services manager **916** with a procedure to follow in the event that events, metadata, or content indicate malicious behavior or code or a violation of any data governance policy. For example, governance policies **936** can contain virus definitions to be used by governance services manager **916** and pattern detection service **934** in order to detect viruses in compromised data objects stored in object store **914**. Upon determining that a data object does contain a virus, governance services manager **916** consults governance policies **936** to determine how to proceed. Governance policies **936** might indicate that the infected data object should be

deleted, quarantined, ignored, etc. Alternatively, governance policies **936** might simply indicate that a notification be sent to an administrator.

[0087] In response to the relevant procedure included in governance policies **936**, governance services manager **916** will begin performing remediation actions. Governance services manager **916** generates control messages, including a list of file system operations to be executed on local file storage system **104** or storage server **106**. Each of the file system operations includes, for example, a data object identifier and one of a set of potential operations, including move, delete, update, etc. Governance service manager **916** saves the control messages into an outgoing queue **938**, which then forwards the messages to one of connector hub interface **902** or cloud connectors **906**, based on the data object identifier. Connector hub interface **902** sends the control messages to local file storage system **104**, which processes the control messages and performs the necessary file system operations. Cloud connectors **906** utilize cloud-based storage APIs to access storage server **106** and perform the necessary file system operations.

[0088] Additionally, data governance services **816(1)** include an administrator interface **940**, which allows the administrator to fine-tune the way in which data governance services **816(1)** detect and respond to data security threats. Administrator interface **940** allows administrators associated with local file storage system **104** or storage server **106** to access event store **910**, metadata store **912**, and/or object store **914** in order to view data indicative of access or changes made to local file storage system **104** and/or storage server **106**. Administrator interface **940** requests the data from governance services manager **916**, which pulls the data and provides it to administrator interface **940**. Administrator interface **940** provides the data to the administrator through a customizable graphical user interface (GUI), which is defined in governance policies **936**. Through the GUI, the administrator can see, for example, what objects are being accessed by who and how frequently, as well as which objects have been changed and in what way. Administrator interface **940** also provides the administrator with options for updating governance policies **936**, including setting remediation procedures, uploading custom content patterns, customizing the GUI, etc.

[0089] FIG. **10A** is a data flow diagram showing a method for initially configuring data governance system **102** and local file storage system **104** for data governance services. First, during a step labeled (1), an administrator **1002** associated with local file storage system **104** requests data governance services from data governance system **102**. In response, during a step labeled (2), data governance system **102** pushes connector hub software **1004** (including any additional agents needed for particular sources) to local file storage system **104**. Upon receiving connector hub software **1004**, local file storage system **104** installs the software onto an appropriate machine (e.g. connector hub host device **406**), during a step labeled (3). Additionally, during a step labeled (4), local file storage system **104** provides a metadata snapshot **1006**, indicative of the file system of local file storage system **104**, to data governance system **102**. Upon receiving metadata snapshot **1006**, data governance system **102** generates a derivative data set indicative of local file storage system **104**, during a step labeled (5).

[0090] FIG. **10B** is a data flow diagram showing a method for initially configuring data governance system **102** and storage server **106** for data governance services. First, during a step labeled (1), administrator **1002**, also associated with storage server **106**, request data governance services from data governance system **102**. In response, data governance system **102** requests a metadata snapshot of storage server **106** (e.g., the files and directories associated with the client), during a step labeled (2). Upon receiving the request, storage server **106** provides a metadata snapshot **1008**, indicative of the file system of the clients data on storage server **106**, during a step labeled (3). Upon receiving metadata snapshot **1008**, data governance system **102** generates a derivative data set indicative of storage server **106**, during a step labeled (4).

[0091] FIG. **10C** is data flow diagram showing a method for updating the derivative data sets stored on data governance system **102**. During a step labeled (1a), changes are made to local file storage system **104**. During another step labeled (1b), changes are also made to storage server **106**.

Responsive to the changes to local file storage system **104**, during a step labeled (2a), local file storage system **104** provides events **1010**, indicative of the changes made on local file storage system **104**, to data governance system **102**. Responsive to the changes to storage server **106**, during another step labeled (2b), data governance system **102** retrieves events **1012**, indicative of the changes to storage server **106**, from storage server **106**. Responsive to receiving events **1010** and **1012**, data governance system **102** updates the derivative data sets indicative of local file storage system **104** and storage server **106**, respectively, during a step labeled (3). It should be noted that local file storage system **104** and storage server **106** can provide additional metadata or content along with events **1010** and **1012**. Additionally, events can be either pushed by local file storage system **104** and storage server **106** or pulled by data governance system **102**, based on the preferences of the client, the configuration of the system, and the capabilities of storage server **106**. [0092] Events **1010** and **1012** are formatted as JSON objects as follows:

```
TABLE-US-00001  {      "messageId" : "<alphanumeric random id>",      "message" : {  
"object" : "<alphanumeric object id>",      "action" : ["create" | "retrieve" | "update" |  
"delete" | ... | ... ],      "username" : "<user_name>",      "actionCreationTime":  
12:34.05_11/31/2017      } }
```

wherein "messageID" is an alphanumeric identifier (e.g., corresponding to an event identifier, a random identifier, etc.), "object" is an alphanumeric identifier that corresponds to a particular object that was accessed/modified, "action" specifies a particular access or modification operation (e.g. a CRUD operation) executed on the particular object, "username" is the username corresponding to the user that executed the operation, and "actionCreationTime" is a timestamp specifying the date and time that the operation was executed on the object. The message can also include, for example, a WebSocket header added to the message (payload) by local file storage system **104** or storage server **106**, if the event was sent via a WebSocket protocol.

[0093] FIG. **11** is a data flow diagram showing a portion of the method of FIG. **10C** in greater detail. Local file storage system **104** includes connector hub **502**, which provides events **1010** to data governance system **102**. Data governance system **102** includes data governance server **806(1)**, which receives events **1010** from local file storage system **104** and events **1012** from storage server **106**. Storage server **106** does not include a connector hub, and data governance server **802(1)** instead utilizes publicly available APIs to request and receive information indicative of events from storage server **106**.

[0094] FIG. **12** is a data flow diagram showing connector hub **502** retrieving events, metadata, and content from sources **1202(1-N)** of data sources **408** and forwarding the events, metadata, and content to data governance system **102**. Each of connectors **602** is associated with a particular one of sources **1202**. For example, connector **602(1)** can be a directory services connector, while source **1202(1)** is a directory service, such as Microsoft Active Directory. When changes are made to sources **1202**, events, content and/or metadata are received from each of sources **1202** by the corresponding one of connectors **602**. Connectors **602** utilize applications **604** to analyze the data pulled from sources **1202**, before utilizing platform services **606** to forward the data to data governance system **102**.

[0095] FIG. **13** is a stacked software diagram of connector hub **502**, including particular ones of connectors **602**, applications **604**, and platform services **606**. Connectors **602** communicate with sources **1202** (FIG. **12**) to capture events, metadata, and content indicative of file system operations executed on sources **1202**, by any combination of registering system callbacks, scanning the sources, utilizing agents on the sources, etc. These methods will be discussed in further detail with reference to FIGS. **14-20**, below. Connector layer **610** includes, by way of non-limiting example, directory service connector **702**, content management connector **704**, messaging connector **706**, file sync and share connector **708**, email connector **710**, and file server connector **712**. Entities of connector layer **610** communicate with entities of application layer **612** and entities of platform services layer **614**. The connectors of connector layer **610** utilize applications **604**, including

content extraction service **714**, content detection service **716**, and pattern detection service **718**, for analysis of the data captured from sources **1202**. Each connector of connector layer **610** and each application of application layer **612** utilize platform services of platform services layer **614**, including source-type connectors **720**, lock manager **722**, distributed configuration **724**, message queue **726**, distributed cache **728**, transport service **730**, and discovery service **732**, to access underlying system hardware and to communicate with one another. Platform services of layer **614** also provide communication with data governance system **102** to provide events, metadata, and content from connectors of connector layer **610**, and any modified or created content from applications of application layer **612**, to data governance system **102**.

[0096] FIGS. **14-20** are process flow diagrams showing example methods for generating and processing events according to the present invention. It should be noted that the particular method used for generating and/or processing events is dependent on the source itself and the preferences of the cloud client. For example, the methods available for generating events on a particular source might be limited to comparing metadata snapshots, or the client may choose to utilize only the most secure method for each source. Other considerations include required processing power/memory usage, time efficiency, etc. Additionally, the various methods are shown originating either in connector hub **502** or data governance system **102**. However, any of the methods shown can originate in either connector hub **502** or data governance system **102**. For example, a method originating in connector hub **502** can originate in data governance system **102**, if data governance system **102** sends a request to connector hub **502** to execute the method.

[0097] FIG. **14** is a process flow diagram showing an example method for retrieving a snapshot of source **1202(N)** via an agent **1402(N)** installed on source **1202(N)**. To initiate the process, connector hub interface **902** sends a “Create_snapshot” request, which includes a request_id to identify the particular request, a source_id to identify source **1202(N)** as the target source, and a volume_id to identify the particular physical or virtual disk(s) of NAS device **402** in which the source is stored. Connector **602(N)** processes the request and, in response, requests a snapshot of the specified disk(s) by sending a “Take_snapshot” request to agent **1402(N)**, including the volume_id. Agent **1402(N)** processes the request and acquires a snapshot of the specified disk(s). Agent **1402(N)** then sends a “snapshot_available” notification, including a snapshot_id, to connector **602(N)**. Connector **602(N)** then forwards the snapshot_available notification to connector hub interface **902**, along with the initial request_id.

[0098] Next, connector hub interface **902** retrieves the folder metadata of the snapshot from connector hub **502**. Connector hub interface **902** sends a “get_folder_snapshot” request to connector hub **502**, including the original request_id, the snapshot_id, the volume_id, and a list of levels to query in the folder tree. Connector hub **502** then walks through the snapshot and gathers the requested folders, one by one. Once the walk is complete, connector hub **502** sends a “snapshot_folder_response”, including the request_id, the snapshot_id, and metadata for the list of requested folder paths. The snapshot_folder_response is sent, for example, via REST APIs in order to facilitate the transfer of large amounts of data over an Internet connection.

[0099] Finally, connector hub interface **902** retrieves the file snapshot from connector hub **502**. In response to receiving the snapshot_folder_response, connector hub interface **902** sends a “get_files_snapshot” request to connector hub **502**, including the request_id, the snapshot_id, and the list of folder paths previously provided by connector hub **502**. Again, connector hub **502** walks through the snapshot and generates a list of the data objects in the specified folders. Connector hub **502** then sends a “file_snapshot_response” to connector hub interface **902**, including the request_id, the snapshot_id, a map of the previously requested folder paths, and a list of file objects contained in each folder of the specified folder paths. The snapshot_file_response is sent, for example, via REST APIs.

[0100] FIGS. **15A** and **15B** illustrate an efficient method for updating the derivative data set located on the cloud, without pulling a full snapshot from the data source located at client site **118**

or cloud-based storage **106** (FIG. 1). Rather than pull a full snapshot of the local file storage system, cloud-based data governance system **102** pulls a folder snapshot (FIG. 15A), updates the folder state based on the folder snapshot, then requests (FIG. 15A) and receives (FIG. 15B) file snapshots only for those folders that were changed.

[0101] FIG. 15A is a process flow diagram showing an example method for processing the folder snapshot in data governance server **806(1)**. Initially, connector hub interface **902** saves the folder snapshot in metadata store **912**, before sending a message, indicating the location of the folder snapshot, to snapshot service **920**, via governance services manager **916**, message processor **908**, and incoming queue **904** (FIG. 9). Snapshot service **920** then downloads the folder snapshot from the specified location on metadata store **912** and utilizes governance services **918** to query a folder state of the derivative data set saved on metadata store **912**. Snapshot service **920** then updates the temporary folder snapshot state and generates file listing operations. Finally, snapshot service **920** sends the batch file listing messages to connector hub interface **902**.

[0102] FIG. 15B is a process flow diagram showing an example method for processing the file snapshot(s) in data governance server **806(1)**. Initially, connector hub interface **902** saves the file snapshot in metadata store **912**, before sending a message, indicating the location of the file snapshot, to snapshot service **920**, via governance services manager **916**, message processor **908**, and incoming queue **904** (FIG. 9). Snapshot service **920** then downloads the file snapshot from the specified location on metadata store **912** and analyzes the snapshot to generate a list of folder update operations. Finally, snapshot service **920** utilizes governance services **918** to update the folders of the derivative data set on metadata store **912** based on the generated folder update operations.

[0103] FIG. 16 is a process flow diagram showing a particular example method for providing a directory service snapshot to data governance server **806(1)**. First, connector hub **502** fetches the directory service snapshot from directory service connector **702**. Then, connector hub **502** pushes the directory service snapshot to connector hub interface **902**, along with the request_id and the source_id (FIG. 14).

[0104] FIG. 17 is a process flow diagram showing an example method for scanning source **1202(N)** to generate events. Initially, connector **602(N)** provides scanning agent software to source **1202(N)**, which installs the scanning agent and returns an installation confirmation. Next, connector **602(N)** sends a scan request to an agent **1702(N)**, which was installed on source **1202(N)**. Agent **1702(N)** performs the scan of source **1202(N)** and returns the scan results. Finally, connector **602(N)** derives events from the scan results and provides the events to connector hub interface **902**.

[0105] FIG. 18 is a process flow diagram showing an example method for intercepting and filtering events via an agent **1802(N)** on the source. Initially, connector **602(N)** sends a request for events to agent **1802(N)**. Agent **1802(N)** queries an event log, generated in response to changes/access to source **1202(N)**, and provides any new events to connector **602(N)**. Then, connector **602(N)** filters the events to remove any redundant/unwanted events. Finally, connector **602(N)** provides the events to connector hub interface **902**.

[0106] FIG. 19 is a process flow diagram showing a particular example method for intercepting and filtering events via agent **1802(N)**. Initially, connector **602(N)** sends a “get_events” request to Agent **1802(N)**, which rolls over the event log and returns the location of a rolled over events file to connector **602(N)**. Next, connector **602(N)** sends a “new_event_packet_available” notification, including a request_id, a source_id, a volume_id, a start_time indicating a time stamp of the first event, and an end_time indicating a time stamp of the final event, to connector hub interface **902**. Finally, connector **602(N)** provides the event packet, including the request_id, to connector hub interface **902**, via REST APIs.

[0107] FIG. 20 is a process flow diagram showing an example method for receiving events from source **1202(N)** by registering for callbacks. Initially, connector **602(N)** registers for callbacks from

source **1202(N)**, whenever source **1202(N)** is accessed or modified. Next, source **1202(N)** is accessed and/or modified by one of local clients **120** (FIG. **1**) and executes the callback by sending events related to the access/modification to connector **602(N)**. Finally, connector **602(N)** filters the events and provides them to connector hub interface **902**.

[0108] FIG. **21** is a flow chart summarizing an example method **2100** of providing data governance services from a site remote from the governed data source. In a first step **2102**, a connection with a remote data storage system associated with a remote file system is established over a wide-area network. Then, in a second step **2104**, a metadata snapshot is received from the remote data storage system. Next, in a third step **2106**, a derivative data set indicative of the remote file system is generated based on the metadata snapshot. Then, in a fourth step **2108**, an event associated with the remote file system is captured. The event is indicative of at least one file system operation executed on a data object of the remote data storage system. Next, in a fifth step **2110**, the derivative data set is updated based on the event. Then, in a sixth step **2112**, data analytics are performed on the derivative data set, after the derivative data set has been updated. Next, in a seventh step **2114**, the event is processed to determine whether the event conflicts with a governance policy of the data governance system. Finally, in an eighth step **2116**, a set of remediation events are executed, if the event conflicts with the governance policy.

[0109] FIG. **22** is a flow chart summarizing an example method of performing fourth step **2108** of method **2100**. In a first step **2202**, an event collection service is deployed to the remote data storage system. Then, in a second step **2204**, file system operations executed on data objects of the remote data storage system are detected with the event collection service. Next, in a third step **2206**, events indicative of the file system operations are generated with the event collection service. Next, in a fourth step **2208**, the events are pushed to the data governance system by the event collection service. Finally, in a fifth step **2210**, the events are received from the remote data storage system via the event collection service.

[0110] FIG. **23** is a flow chart summarizing an example method **2300** of utilizing data governance services from a remote site. In a first step **2302**, a metadata snapshot indicative of a local data storage system is generated. Then, in a second step **2304**, a connection is established with a remote cloud-based data governance system over a wide-area network. Next, in a third step **2306**, the metadata snapshot is provided to the data governance system. Then, in a fourth step **2308**, an event indicative of a file system operation performed on a data object stored in the local data storage system is captured. Next, in a fifth step **2310**, the event is provided to the data governance system to facilitate a determination of whether the event conflicts with a data governance policy stored on the data governance system. Finally, in a sixth step **2312**, responsive to one or more communication from the data governance system, a set of remediation actions is executed on the local data storage system, if the data governance system determines that the event conflicted with the data governance policy.

[0111] FIG. **24** is a flow chart summarizing an example method of performing fourth step **2308** and fifth step **2310** of method **2300**. In a first step **2402**, an event collection service is received from the data governance system. The event collection service includes a plurality of data monitors that are each associated with one of a plurality of different data source types. The data monitors are operative to detect file system operations executed on data objects of the associated data source type. Next, in a second step **2404**, the event collection service is deployed on the local data storage system. Then, in a third step **2406**, events indicative of file system operations on the local data storage system are generated by the event collection service. Finally, in a fourth step **2408**, the generated events are pushed to the remote data governance system from the event collection service.

[0112] The description of particular embodiments of the present invention is now complete. Many of the described features may be substituted, altered or omitted without departing from the scope of the invention. For example, alternate connectors (e.g., smartphone connectors, CAD connectors,

etc.), may augment or be substituted for any of the example connectors **602**. As another example, additional event collection mechanisms can be used by connector hub **502** or cloud connectors **906**, in order to generate file system events. These and other deviations from the particular embodiments shown will be apparent to those skilled in the art, particularly in view of the foregoing disclosure.

Claims

- 1.** In a cloud-based data governance system, a method for providing data governance of a remote data storage system associated with a remote file system, said method comprising: establishing a connection with said remote data storage system over a wide area network (WAN); capturing an event associated with said remote file system, said event being indicative of at least one file system operation executed on a data object of said remote data storage system; processing said event to determine whether said event conflicts with a governance policy of said data governance system; and executing a set of remediation actions, if said event does conflict with said governance policy.
- 2.** The method of claim 1, wherein said step of capturing an event associated with said remote file system includes: deploying an event collection service to said remote data storage system, said event collection service being operative to detect file system operations executed on data objects of said remote data storage system, generate events indicative of said file system operations, and push said events to said data governance system; and receiving said events from said remote data storage system via said event collection service.
- 3.** The method of claim 1, further comprising: receiving a metadata snapshot from said remote data storage system, said metadata snapshot being indicative of said remote file system; and generating a derivative data set indicative of said remote file system based on said metadata snapshot.
- 4.** The method of claim 3, wherein said step of capturing an event associated with said remote file system includes capturing metadata associated with one or both of said at least one file system operation and said data object.
- 5.** The method of claim 4, wherein said step of capturing metadata includes capturing metadata indicative of a particular user executing said at least one file system operation.
- 6.** The method of claim 5, wherein said step of executing a set of remediation actions includes altering permissions associated with said particular user.
- 7.** The method of claim 3, wherein said step of processing said event includes: updating said derivative data set based on said event; and performing data analytics on said derivative data set after said derivative data set has been updated.
- 8.** The method of claim 1, wherein said step of executing a set of remediation actions includes pushing a control message to said remote data storage system, said control message indicating a set of file system operations to be executed on objects of said remote file system by said remote data storage system.
- 9.** The method of claim 1, further comprising: collecting additional events, each event of said additional events being indicative of at least one additional file system operation executed on a data object of said remote file system stored on said remote data storage system; storing said event and said additional events in an event database; and providing a client associated with said remote file storage system access to said event database.
- 10.** The method of claim 1, wherein said step of processing said event includes performing data analytics on said event.
- 11.** The method of claim 1, wherein said step of establishing a connection with said remote data storage system includes establishing a connection with a third party cloud service provider.
- 12.** A cloud-based data governance system comprising: a processing unit configured to execute code; a network adapter electrically coupled to establish a connection with a remote data storage system associated with a remote file system over a wide-area network (WAN); and memory for storing data and said code, said data and said code including an event collection interface

configured to capture an event from said remote data storage system, said event being indicative of at least one file system operation executed on a data object of said remote file system stored on said remote data storage system, a data governance service configured to receive said event from said event collection interface and to process said event to determine whether said at least one file system operation conflicts with a governance policy of said data governance system, and an enforcement service configured to execute a set of remediation actions, if said at least one file system operation does conflict with said governance policy.

13. The system of claim 12, wherein said event collection interface is configured to: deploy an event collection service to said remote file storage system, said event collection service being operative to detect file system operations executed on data objects of said remote file system stored on said remote data storage system, generate events indicative of said file system operations, and push said events to said data governance system; and said event collection interface is configured to receive said events from said remote data storage system via said event collection service.

14. The system of claim 12, wherein: said event collection interface is further configured to receive a metadata snapshot of said remote data storage system, said metadata snapshot being indicative of said remote file system; and said data governance service is further configured to generate a derivative data set based on said metadata snapshot.

15. The system of claim 14, wherein said event collection interface is configured to capture metadata associated with one or both of said at least one file system operation and said data object associated with said event.

16. The system of claim 15, wherein said event collection interface is configured to capture metadata indicative of a particular user executing said at least one file system operation.

17. The system of claim 16, wherein said set of remediation actions includes altering permissions associated with said particular user.

18. The system of claim 14, wherein said data governance service is additionally configured to update said derivative data set based on said event and to perform data analytics on said updated derivative data set.

19. The system of claim 12, wherein said enforcement module is additionally configured to push one or more control messages to said remote data storage system, said control message(s) indicating a set of file system operations to be executed on objects of said file system on said remote data storage system.

20. The system of claim 12, further comprising: an event database operative to store a record of said event; and a client interface configured to provide a client associated with said remote file system access to said event database; and wherein said event collection interface is configured to collect additional events and store records of said additional events in said database, each event of said additional events being indicative of at least one additional file system operation executed on a data object of said remote file system stored on said remote data storage system; and said data governance service is additionally configured to perform batch data analysis functions on a subset of said records of said database.

21. The system of claim 12, wherein said data governance service is additionally configured to perform data analytics on said event.

22. The system of claim 12, wherein said remote computer system is a third party cloud service provider.
