

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265464

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Power; Martin et al.

METHODS AND APPARATUS TO PERFORM MACHINE-LEARNING MODEL OPERATIONS ON SPARSE ACCELERATORS

Abstract

Methods, apparatus, systems and articles of manufacture are disclosed to perform machine-learning model operations on sparse accelerators. An example apparatus includes first circuitry, second circuitry to generate sparsity data based on an acceleration operation, and third circuitry to instruct one or more data buffers to provide at least one of activation data or weight data based on the sparsity data to the first circuitry, the first circuitry to execute the acceleration operation based on the at least one of the activation data or the weight data.

Inventors: Power; Martin (Dublin, IE), Brady; Kevin (Newry, GB), Hanrahan; Niall (Galway, IE), Grymel; Martin-Thomas (Norderstedt, DE), Bernard; David (Kilcullen, IE), Baugh; Gary (Bray, IE)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 1000008578183

Assignee: Intel Corporation (Santa Clara, CA)

Appl. No.: 19/198452

Filed: May 05, 2025

Related U.S. Application Data

parent US continuation 17357924 20210624 parent-grant-document US 12321857 child US 19198452

Publication Classification

Int. Cl.: G06N3/08 (20230101); G06F17/15 (20060101); G06N3/063 (20230101)

Background/Summary

CROSS REFERENCE TO RELATED APPLICATION [0001] This application is a continuation of (and claims priority to) U.S. patent application Ser. No. 17/357,924, filed Jun. 24, 2021, titled “METHODS AND APPARATUS TO PERFORM MACHINE-LEARNING MODEL OPERATIONS ON SPARSE ACCELERATORS,” which is incorporated herein by reference in its entirety for all purposes.

FIELD OF THE DISCLOSURE

[0002] This disclosure relates generally to processors and, more particularly, to methods and apparatus to perform machine-learning model operations on sparse accelerators.

BACKGROUND

[0003] In recent years, a demand for computationally-intensive processing capabilities, such as image processing capabilities, has moved beyond high-power dedicated desktop hardware and has become an expectation for personal and/or otherwise mobile devices. Mobile devices typically include processing capabilities that are limited by size constraints, temperature management constraints, and/or supply power constraints.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 is an illustration of an example computing system including an example accelerator compiler to configure example acceleration circuitry based on an acceleration operation to be executed by the acceleration circuitry.

[0005] FIG. 2 is an illustration of an example implementation of the example acceleration circuitry of FIG. 1.

[0006] FIG. 3 is a block diagram of an example implementation of the example accelerator compiler of FIG. 1.

[0007] FIG. 4 is an illustration of an example conventional convolution operation that may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0008] FIG. 5 is an illustration of an example depthwise convolution operation that may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0009] FIG. 6 is an illustration of example sparsity bit mask vectors that may be implemented by the example acceleration circuitry of FIGS. 1 and/or 2.

[0010] FIG. 7 is an illustration of an example grouped convolution operation that may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0011] FIG. 8 is an illustration of example bit masks that may be utilized to implement the example grouped convolution operation of FIG. 7, which may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0012] FIG. 9 is an illustration of an example elementwise addition operation that may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0013] FIG. 10 is an illustration of another example elementwise addition operation that may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0014] FIG. 11 is an illustration of yet another example elementwise addition operation that may be

executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0015] FIG. 12 is an illustration of an example dilated convolution operation that may be executed by the example acceleration circuitry of FIGS. 1 and/or 2.

[0016] FIG. 13 is an illustration of example kernels that may be implemented by the example acceleration circuitry of FIGS. 1 and/or 2.

[0017] FIG. 14 is a flowchart representative of an example process that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the example acceleration circuitry of FIGS. 1 and/or 2 to execute an acceleration operation.

[0018] FIG. 15 is another flowchart representative of an example process that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the example acceleration circuitry of FIGS. 1 and/or 2 to execute an acceleration operation.

[0019] FIG. 16 is a flowchart representative of an example process that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the example acceleration circuitry of FIGS. 1 and/or 2 to execute an acceleration operation based on an acceleration mode.

[0020] FIG. 17 is a flowchart representative of an example process that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the example accelerator compiler of FIGS. 1 and/or 3 to generate an executable based on a compiled configuration of a machine-learning model.

[0021] FIG. 18 is a block diagram of an example processing platform structured to execute the example machine readable instructions of FIGS. 15, 16, and/or 17 to implement the example accelerator compiler of FIGS. 1 and/or 3 and/or the example acceleration circuitry of FIGS. 1 and/or 2.

[0022] FIG. 19 is a block diagram of an example implementation of the processor circuitry of FIG. 18.

[0023] FIG. 20 is a block diagram of another example implementation of the processor circuitry of FIG. 18.

[0024] FIG. 21 is a block diagram of an example software distribution platform to distribute software to client devices.

DETAILED DESCRIPTION

[0025] The figures are not to scale. In general, the same reference numbers will be used throughout the drawing(s) and accompanying written description to refer to the same or like parts. Connection references (e.g., attached, coupled, connected, and joined) are to be construed broadly and may include intermediate members between a collection of elements and relative movement between elements unless otherwise indicated. As such, connection references do not necessarily infer that two elements are directly connected and in fixed relation to each other.

[0026] Descriptors “first,” “second,” “third,” etc., are used herein when identifying multiple elements or components which may be referred to separately. Unless otherwise specified or understood based on their context of use, such descriptors are not intended to impute any meaning of priority, physical order or arrangement in a list, or ordering in time but are merely used as labels for referring to multiple elements or components separately for ease of understanding the disclosed examples. In some examples, the descriptor “first” may be used to refer to an element in the detailed description, while the same element may be referred to in a claim with a different descriptor such as “second” or “third.” In such instances, it should be understood that such descriptors are used merely for ease of referencing multiple elements or components.

[0027] As used herein, “processor circuitry” is defined to include (i) one or more special purpose electrical circuits structured to perform specific operation(s) and including one or more semiconductor-based logic devices (e.g., electrical hardware implemented by one or more

transistors), and/or (ii) one or more general purpose semiconductor-based electrical circuits programmed with instructions to perform specific operations and including one or more semiconductor-based logic devices (e.g., electrical hardware implemented by one or more transistors). Examples of processor circuitry include programmed microprocessors, Field Programmable Gate Arrays (FPGAs) that may instantiate instructions, Central Processor Units (CPUs), Graphics Processor Units (GPUs), Digital Signal Processors (DSPs), XPU, or microcontrollers and integrated circuits such as Application Specific Integrated Circuits (ASICs). [0028] For example, an XPU may be implemented by a heterogeneous computing system including multiple types of processor circuitry (e.g., one or more FPGAs, one or more CPUs, one or more GPUs, one or more DSPs, etc., and/or a combination thereof) and application programming interface(s) (API(s)) that may assign computing task(s) to whichever one(s) of the multiple types of the processing circuitry is/are best suited to execute the computing task(s).

[0029] Typical computing systems, including personal computers and/or mobile devices, implement computationally-intensive tasks, such as advanced image processing or computer vision algorithms to automate tasks that human vision can perform. For example, computer vision tasks may include acquiring, processing, analyzing, and/or understanding digital images. Some such tasks facilitate, in part, extraction of dimensional data from the digital images to produce numerical and/or symbolic information. Computer vision algorithms can use the numerical and/or symbolic information to make decisions and/or otherwise perform operations associated with three-dimensional (3-D) pose estimation, event detection, object recognition, video tracking, etc., among others. To support augmented reality (AR), virtual reality (VR), robotics, and/or other applications, it is then accordingly important to perform such tasks quickly (e.g., substantially in real time or near real time) and efficiently.

[0030] Computationally-intensive tasks, such as advanced image processing or computer vision algorithms, may be implemented utilizing a neural network, such as a convolutional neural network (CNN, or ConvNet). A CNN is a deep, artificial neural network (ANN) typically used to classify images, cluster the images by similarity (e.g., a photo search), and/or perform object recognition within the images using convolution. As used herein, “convolution” is defined to be a function derived from two given functions by integration that expresses how a shape of one of the functions is modified by a shape of the other function. Thus, a CNN can be used to identify faces, individuals, street signs, animals, etc., included in an input image by passing an output of one or more filters corresponding to an image feature (e.g., a horizontal line, a two-dimensional (2-D) shape, etc.) over the input image to identify matches of the image feature within the input image.

[0031] CNNs obtain vectors (e.g., broken down from multi-dimensional arrays) that need to be stored or used in computations to perform one or more functions. Thus, a CNN may receive multi-dimensional arrays (e.g., tensors or rows of vectors) including data corresponding to one or more images. The multi-dimensional arrays are broken into vectors. Such vectors may include thousands of elements. Each such element may include a large number of bits. A vector with 10,000 16-bit elements corresponds to 160,000 bits of information. Storing such vectors requires a lot of memory. However, such vectors may include large numbers of elements with a value of zero. Accordingly, some CNNs or other processing engines may break up such a vector into a sparse vector and a sparsity map vector (e.g., a bitmap vector).

[0032] As defined herein, a “sparse vector” is a vector that includes all non-zero elements of a vector in the same order as a dense vector, but excludes all zero elements. As defined herein, a “dense vector” is an input vector including both zero and non-zero elements. As such, the dense vector [0, 0, 5, 0, 18, 0, 4, 0] corresponds to the sparse vector is [5, 18, 4]. As defined herein, a “sparsity map” is a vector that includes one-bit elements identifying whether respective elements of the dense vector is zero or non-zero. Thus, a sparsity map may map non-zero values of the dense vector to ‘1’ and may map the zero values of the dense vector to ‘0’. For the above-dense vector of [0, 0, 5, 0, 18, 0, 4, 0], the sparsity map may be [0, 0, 1, 0, 1, 0, 1, 0] (e.g., because the third, fifth,

and seventh elements of the dense vector are non-zero). The combination of the sparse vector and the sparsity map represents the dense vector (e.g., the dense vector could be generated/reconstructed based on the corresponding sparse vector and sparsity map). Accordingly, a CNN engine may generate/determine the dense vector based on the corresponding sparse vector and sparsity map without storing the dense vector in memory.

[0033] Storing a sparse vector and a sparsity map in memory instead of a dense vector saves memory and processing resources (e.g., providing there are sufficient zeros in the dense vector(s)). For example, if each element of the above-dense vector (e.g., [0, 0, 5, 0, 18, 0, 4, 0]) was 16 bits of information, the amount of memory required to store the dense vector is 128 bits (e.g., $128 \text{ bits} = 8 \text{ elements} \times 16 \text{ bits}$). However, the amount of memory required to store the corresponding sparse vector (e.g., [5, 18, 4]) and the sparsity map (e.g., [0, 0, 1, 0, 1, 0, 1, 0]) is 64 bits (e.g., $(3 \text{ elements of the sparse vector} \times 16 \text{ bits}) + (8 \text{ elements of the sparsity map} \times 1 \text{ bit})$). Accordingly, storing the sparse vector and sparsity map instead of a corresponding dense vector reduces the amount of memory needed to store such vectors. Additionally, utilizing sparse vectors and sparsity maps improves bandwidth requirements because a decrease of an amount of data being delivered into a computational engine increases a delivery speed to the computational engine.

[0034] Machine-learning accelerators, such as neural network accelerators (NNAs), may execute 2-D convolution operations with dedicated hardware. Architectures of machine-learning models that may be executed by the machine-learning accelerators are continuously evolving. Some such architectures incorporate machine-learning layers that are not conventional 2-D convolution operations. For example, non-2-D convolution operations may include depthwise convolution, elementwise addition, etc. Machine-learning accelerators tailored to facilitate execution of 2-D convolution may be unable to execute such non-2-D convolutions.

[0035] Some machine-learning accelerators may support such non-2-D convolution operations by offloading the non-2-D convolution operations to other hardware. For example, a 2-D convolution machine-learning accelerator may offload layers to different hardware to execute depthwise convolution layers and retrieve the outputs from the different hardware to facilitate another machine-learning task. In some such examples, the offloading of the layers and retrieval of the outputs increases processing latency when executing non-2-D convolution operations. For example, the machine-learning accelerator may become idle as the layer is offloaded, which may substantially degrade the performance of the machine-learning accelerator. In some such examples, the performance of the machine-learning accelerator may also be degraded due to traversing a memory hierarchy to supply the offloaded layers as inputs to the different hardware and to traverse the memory hierarchy again for the machine-learning accelerator to receive the output from the different hardware and continue processing the machine-learning network.

[0036] Some machine-learning accelerators may support non-2-D convolution operations with dedicated hardware. For example, a machine-learning accelerator may include first hardware to execute a 2-D convolution operation and second hardware to execute a non-2-D convolution operation. In some such examples, the inclusion of dedicated hardware within a machine-learning accelerator to execute the non-2-D convolution operations increases the overall size and/or power consumption of the machine-learning accelerator.

[0037] In some instances, adding dedicated hardware for layers such as depthwise and elementwise addition may directly degrade the performance of the machine-learning accelerator. For example, performance of a machine-learning accelerator may be measured in Tera-operations per second (TOPs) and the efficiency of the machine-learning accelerator may be measured in TOPs per square millimeter of die area (TOPs/mm²) and/or TOPs per watt of device power (TOPs/W). Accordingly, the inclusion of dedicated hardware may directly degrade TOPs/mm² of the machine-learning accelerator because increased die area to execute the same number of operations may reduce the TOPs/mm² metric. The inclusion of dedicated hardware may also require the inclusion of additional data paths within the machine-learning accelerator and multiplexing between these data

paths adds significant complexity to the physical implementation of the machine-learning accelerator when targeting substantial frequencies of operation (e.g., gigahertz (GHz) frequencies of operation).

[0038] Examples disclosed herein perform machine-learning model operations on sparse accelerators. As used herein, a “sparse accelerator” is a hardware accelerator that may implement machine-learning operations (e.g., convolution operations) by utilizing sparse data structures (e.g., sparse vectors, sparsity map vectors, etc.). In some disclosed examples, an example sparse accelerator may generate sparsity information internally within the sparse accelerator. For example, the sparse accelerator may generate sparse or sparsity information including activation sparsity data (e.g., sparse activation data, sparse activation vector(s), etc.) and/or weight sparsity data (e.g., sparse weight data, sparse weight vector(s), etc.). In some such examples, the sparse accelerator may achieve increased bandwidth by internally generating the sparsity information when compared to retrieving the sparsity information from external hardware. In some disclosed examples, the sparse accelerator may generate weight data internally within the sparse accelerator. In some such disclosed examples, the sparse accelerator may achieve increased bandwidth by internally generating the weight information when compared to retrieving the weight information from external hardware.

[0039] In some disclosed examples, the sparse accelerator may adjust and/or otherwise manipulate sparsity bit masks to implement non-2-D convolution operations on the sparse accelerator, which may include depthwise convolution, grouped convolution, non-convolution operations (e.g., elementwise addition), dilated convolution, etc. In some disclosed examples, the sparse accelerator may implement a custom operation in which a user, a computing system, etc., may identify whether at least one of internal weight data generation, internal weight sparsity generation, or internal activation sparsity generation is to be executed by the sparse accelerator.

[0040] Advantageously, the example sparse accelerator disclosed herein may execute non-2-D convolution layers as 2-D convolution layers on a sparse machine-learning architecture by invoking sparsity bit mask manipulation, internal generation of sparsity information, internal generation of weight data, etc., and/or a combination thereof. Advantageously, the example sparse accelerator disclosed herein may achieve improved power and performance compared to prior accelerators by executing non-convolution operations as convolutions operations. Advantageously, the example sparse accelerator disclosed herein may achieve increased efficiencies of executing machine-learning networks (e.g., neural networks) on sparse machine-learning accelerators (e.g., sparse neural networks (NNs)) by executing the non-convolution operations as convolution operations.

[0041] In some disclosed examples, the sparse accelerator may program configuration registers within the sparse accelerator to internally generate bit mask vectors, weight data, etc., to manipulate the values that are subject to computation and implement the layer operation(s) as a convolution. Advantageously, internal generation of the bit mask vectors, the weight data, etc., may save the bandwidth of reading bit mask vectors, the weight data, etc., from external resources. In some disclosed examples, the sparse accelerator may manipulate the bit mask vectors values in memory and read in these vectors to manipulate the values that are subject to computation and implement the layer operation(s) as a convolution.

[0042] FIG. 1 is an illustration of an example computing environment **100** including an example computing system **102**, which includes an example accelerator compiler **104A-C** to configure a sparse accelerator to execute non-convolution operations as convolution operations to achieve improved accelerator efficiency and performance. In some examples, the accelerator compiler **104A-C** obtains an output from a machine-learning framework (e.g., a NN framework) and compiles the output for implementation on the sparse accelerator based on the non-convolution operation to be executed and/or otherwise performed by the sparse accelerator.

[0043] The computing system **102** of the illustrated example of FIG. 1 includes an example central processing unit (CPU) **106**, a first example acceleration circuitry (ACCELERATION CIRCUITRY

A) **108**, a second example acceleration circuitry (ACCELERATION CIRCUITRY B) **110**, an example general purpose processing circuitry **112**, an example interface circuitry **114**, an example bus **116**, an example power source **118**, and an example datastore **120**. In this example, the datastore **120** includes example configuration data (CONFIG DATA) **122** and example machine-learning model(s) (ML MODEL(S)) **124**. Further depicted in the illustrated example of FIG. **1** are an example user interface **126**, an example network **128**, and example external computing systems **130**.

[0044] In some examples, the computing system **102** is a system on a chip (SoC) representative of one or more integrated circuits (ICs) (e.g., compact ICs) that incorporate components of a computer or other electronic system in a compact format. For example, the computing system **102** may be implemented with a combination of one or more programmable processors, hardware logic, and/or hardware peripherals and/or interfaces. Additionally or alternatively, the example computing system **102** of FIG. **1** may include memory, input/output (I/O) port(s), and/or secondary storage. For example, the computing system **102** includes the acceleration compiler **104A-C**, the CPU **106**, the first acceleration circuitry **108**, the second acceleration circuitry **110**, the general purpose processing circuitry **112**, the interface circuitry **114**, the bus **116**, the power source **118**, the datastore **120**, the memory, the I/O port(s), and/or the secondary storage all on the same substrate (e.g., silicon substrate, semiconductor-based substrate, etc.). In some examples, the computing system **102** includes digital, analog, mixed-signal, radio frequency (RF), or other signal processing functions.

[0045] In the illustrated example of FIG. **1**, the first acceleration circuitry **108** is an artificial intelligence (AI) accelerator. For example, the first acceleration circuitry **108** may be implemented by a hardware accelerator configured to accelerate AI tasks or workloads, such as NNs (e.g., artificial neural networks (ANNs)), machine vision, machine learning, etc. In some examples, the first acceleration circuitry **108** may implement a sparse accelerator (e.g., a sparse hardware accelerator). In some examples, the first acceleration circuitry **108** may implement a vision processing unit (VPU) to effectuate machine or computer vision computing tasks, train and/or execute a physical neural network, and/or train and/or execute a neural network. In some examples, the first acceleration circuitry **108** may train and/or execute a convolution neural network (CNN), a deep neural network (DNN), an ANN, a recurrent neural network (RNN), etc., and/or a combination thereof.

[0046] In the illustrated example of FIG. **1**, the second acceleration circuitry **110** is a graphics processing unit (GPU). For example, the second acceleration circuitry **110** may be a GPU that generates computer graphics, executes general-purpose computing, etc. In some examples, the second acceleration circuitry **110** is another instance of the first acceleration circuitry **108**. In some such examples, the computing system **102** may provide portion(s) of AI/ML workloads to be executed in parallel by the first acceleration circuitry **108** and the second acceleration circuitry **110**.

[0047] The general purpose processing circuitry **112** of the example of FIG. **1** is a programmable processor, such as a CPU or a GPU. Alternatively, one or more of the first acceleration circuitry **108**, the second acceleration circuitry **110**, and/or the general purpose processing circuitry **112** may be a different type of hardware such as a digital signal processor (DSP), an application specific integrated circuit (ASIC), a programmable logic device (PLD), and/or a field programmable logic device (FPLD) (e.g., a field-programmable gate array (FPGA)).

[0048] In the illustrated example of FIG. **1**, the interface circuitry **114** is hardware that may implement one or more interfaces (e.g., computing interfaces, network interfaces, etc.). For example, the interface circuitry **114** may be hardware, software, and/or firmware that implements a communication device (e.g., a network interface card (NIC), a smart NIC, a gateway, a switch, etc.) such as a transmitter, a receiver, a transceiver, a modem, a residential gateway, a wireless access point, and/or a network interface to facilitate exchange of data with external machines (e.g., computing devices of any kind) via the network **128**. In some examples, the communication is

effectuated via a Bluetooth® connection, an Ethernet connection, a digital subscriber line (DSL) connection, a wireless fidelity (Wi-Fi) connection, a telephone line connection, a coaxial cable system, a satellite system, a line-of-site wireless system, a cellular telephone system, an optical connection (e.g., a fiber-optic connection), etc. For example, the interface circuitry **114** may be implemented by any type of interface standard, such as a Bluetooth® interface, an Ethernet interface, a Wi-Fi interface, a universal serial bus (USB), a near field communication (NFC) interface, and/or a peripheral component interconnect express (PCIe) interface.

[0049] The computing system **102** includes the power source **118** to deliver power to hardware of the computing system **102**. In some examples, the power source **118** may implement a power delivery network. For example, the power source **118** may implement an alternating current-to-direct current (AC/DC) power supply. In some examples, the power source **118** may be coupled to a power grid infrastructure such as an AC main (e.g., a 110 volt (V) AC grid main, a 220 V AC grid main, etc.). Additionally or alternatively, the power source **118** may be implemented by a battery. For example, the power source **118** may be a limited energy device, such as a lithium-ion battery or any other chargeable battery or power source. In some such examples, the power source **118** may be chargeable using a power adapter or converter (e.g., an AC/DC power converter), a wall outlet (e.g., a 110 V AC wall outlet, a 220 V AC wall outlet, etc.), a portable energy storage device (e.g., a portable power bank, a portable power cell, etc.), etc.

[0050] The computing system **102** of the illustrated example of FIG. **1** includes the datastore **120** to record data (e.g., the configuration data **122**, the ML model(s) **124**, etc.). The datastore **120** of this example may be implemented by a volatile memory (e.g., a Synchronous Dynamic Random Access Memory (SDRAM), a Dynamic Random Access Memory (DRAM), a RAMBUS Dynamic Random Access Memory (RDRAM), etc.) and/or a non-volatile memory (e.g., flash memory). The datastore **120** may additionally or alternatively be implemented by one or more double data rate (DDR) memories, such as DDR, DDR2, DDR3, DDR4, mobile DDR (mDDR), etc. The datastore **120** may additionally or alternatively be implemented by one or more mass storage devices such as hard disk drive(s) (HDD(s)), compact disk (CD) drive(s), digital versatile disk (DVD) drive(s), solid-state disk (SSD) drive(s), etc. While in the illustrated example the datastore **120** is illustrated as a single datastore, the datastore **120** may be implemented by any number and/or type(s) of datastores. Furthermore, the data stored in the datastore **120** may be in any data format such as, for example, binary data, comma delimited data, tab delimited data, structured query language (SQL) structures, an executable, etc.

[0051] In the illustrated example of FIG. **1**, the computing system **102** is in communication with the user interface **126**. For example, the user interface **126** may be implemented by a graphical user interface (GUI), an application user interface, etc., which may be presented to a user on a display device in circuit with and/or otherwise in communication with the computing system **102**. In some such examples, a user (e.g., a developer, an IT administrator, a customer, etc.) controls the computing system **102**, configures, trains, and/or executes the ML model(s) **124**, etc., via the user interface **126**. Alternatively, the computing system **102** may include and/or otherwise implement the user interface **126**.

[0052] In the illustrated example of FIG. **1**, the accelerator compiler **104A-C**, the CPU **106**, the first acceleration circuitry **108**, the second acceleration circuitry **110**, the general purpose processing circuitry **112**, the interface circuitry **114**, the power source **118**, and the datastore **120** are in communication with one(s) of each other via the bus **116**. For example, the bus **116** may be implemented by at least one of an Inter-Integrated Circuit (I2C) bus, a Serial Peripheral Interface (SPI) bus, a Peripheral Component Interconnect (PCI) bus, or a PCIe bus. Additionally or alternatively, the bus **116** may be implemented by any other type of computing or electrical bus.

[0053] In the illustrated example of FIG. **1**, the network **128** is the Internet. However, the network **128** of this example may be implemented using any suitable wired and/or wireless network(s) including, for example, one or more data buses, one or more Local Area Networks (LANs), one or

more wireless LANs, one or more cellular networks, one or more private networks, one or more public networks, etc. In some examples, the network **128** enables the computing system **102** to be in communication with one(s) of the external computing systems **130**.

[0054] In the illustrated example of FIG. **1**, the external computing systems **130** include and/or otherwise implement one or more computing devices on which the ML model(s) **124** is/are to be executed. In this example, the external computing systems **130** include an example desktop computer **132**, an example mobile device (e.g., a smartphone, an Internet-enabled smartphone, etc.) **134**, an example laptop computer **136**, an example tablet (e.g., a tablet computer, an Internet-enabled tablet computer, etc.) **138**, and an example server **140**. In some examples, fewer or more than the external computing systems **130** depicted in FIG. **1** may be used. Additionally or alternatively, the external computing systems **130** may include, correspond to, and/or otherwise be representative of, any other type and/or quantity of computing devices.

[0055] In some examples, one or more of the external computing systems **130** execute one(s) of the ML model(s) **124** to process a computing workload (e.g., an AI/ML workload). For example, the mobile device **134** can be implemented as a cell or mobile phone having one or more processors (e.g., a CPU, a GPU, a VPU, an AI or NN specific processor, etc.) on a single SoC to process an AI/ML workload using one(s) of the ML model(s) **124**. In some examples, the desktop computer **132**, the laptop computer **136**, the tablet computer, and/or the server **140** may be implemented as computing device(s) having one or more processors (e.g., a CPU, a GPU, a VPU, an AI/NN specific processor, etc.) on one or more SoCs to process AI/ML workload(s) using one(s) of the ML model(s) **124**. In some examples, the server **140** may implement one or more servers (e.g., physical servers, virtualized servers, etc., and/or a combination thereof) that may implement a data facility, a cloud service (e.g., a public or private cloud provider, a cloud-based repository, etc.), etc., to process AI/ML workload(s) using one(s) of the ML model(s) **124**.

[0056] In the illustrated example of FIG. **1**, the computing system **102** includes a first accelerator compiler **104A** (e.g., a first instance of the accelerator compiler **104A-C**), a second accelerator compiler **104B** (e.g., a second instance of the accelerator compiler **104A-C**), and a third accelerator compiler **104C** (e.g., a third instance of the accelerator compiler **104A-C**) (collectively referred to herein as the accelerator compiler **104A-C** unless specified otherwise). In this example, the first accelerator compiler **104A** is implemented by the CPU **106** (e.g., implemented by hardware, software, and/or firmware of the CPU **106**).

[0057] In the illustrated example of FIG. **1**, the second accelerator compiler **104B** is implemented by the general purpose processing circuitry **112** (e.g., implemented by hardware, software, and/or firmware of the general purpose processing circuitry **112**). In this example, the third accelerator compiler **104C** is external to the CPU **106**. For example, the third accelerator compiler **104C** may be implemented by hardware, software, and/or firmware of the computing system **102**. In some such examples, the third accelerator compiler **104C** may be implemented by one or more analog or digital circuit(s), logic circuits, programmable processor(s), programmable controller(s), GPU(s), DSP(s), ASIC(s), PLD(s), and/or FPLD(s)).

[0058] In some examples, one or more of the first accelerator compiler **104A**, the second accelerator compiler **104B**, the third accelerator compiler **104C**, and/or portion(s) thereof, may be virtualized, such as by being implemented with one or more containers, one or more virtual resources (e.g., virtualizations of compute, memory, networking, storage, etc., physical hardware resources), one or more virtual machines, etc. In some examples one or more of the first accelerator compiler **104A**, the second accelerator compiler **104B**, the third accelerator compiler **104C**, and/or portion(s) thereof, may be implemented by different resource(s) of the computing system **102**.

Alternatively, the computing system **102** may not include one or more of the first accelerator compiler **104A**, the second accelerator compiler **104B**, and/or the third accelerator compiler **104C**.

[0059] In the illustrated example of FIG. **1**, the accelerator compiler **104A-C** may compile an AI/ML framework based on the configuration data **122** for implementation on one(s) of the

acceleration circuitry **108, 110**. In some examples, the configuration data **122** may include AI/ML configuration data (e.g., register configurations, activation data, activation sparsity data, weight data, weight sparsity data, hyperparameters, etc.), a convolution operation to be executed (e.g., a 2-D convolution, a depthwise convolution, a grouped convolution, a dilated convolution, etc.), a non-convolution operation (e.g., an elementwise addition operation), etc., and/or a combination thereof. In some examples, the accelerator compiler **104A-C** may compile the AI/ML framework to generate an executable construct that may be executed by the one(s) of the acceleration circuitry **108, 110**.

[0060] In the illustrated example of FIG. 1, the accelerator compiler **104A-C** may instruct, direct, and/or otherwise invoke one(s) of the acceleration circuitrys **108, 110** to execute one(s) of the ML model(s) **124**. For example, the ML model(s) **124** may implement AI/ML models. AI, including machine learning (ML), deep learning (DL), and/or other artificial machine-driven logic, enables machines (e.g., computers, logic circuits, etc.) to use a model to process input data to generate an output based on patterns and/or associations previously learned by the model via a training process. For instance, the machine-learning model(s) **124** may be trained with data to recognize patterns and/or associations and follow such patterns and/or associations when processing input data such that other input(s) result in output(s) consistent with the recognized patterns and/or associations.

[0061] Many different types of machine-learning models and/or machine-learning architectures exist. In some examples, the accelerator compiler **104A-C** generates the machine-learning model(s) **124** as neural network model(s). The accelerator compiler **104A-C** may invoke the interface circuitry **114** to transmit the machine-learning model(s) **124** to one(s) of the external computing systems **130**. Using a neural network model enables the acceleration circuitry **108, 110** to execute an AI/ML workload. In general, machine-learning models/architectures that are suitable to use in the example approaches disclosed herein include recurrent neural networks. However, other types of machine learning models could additionally or alternatively be used such as supervised learning ANN models, clustering models, classification models, etc., and/or a combination thereof. Example supervised learning ANN models may include two-layer (2-layer) radial basis neural networks (RBN), learning vector quantization (LVQ) classification neural networks, etc. Example clustering models may include k-means clustering, hierarchical clustering, mean shift clustering, density-based clustering, etc. Example classification models may include logistic regression, support-vector machine or network, Naive Bayes, etc. In some examples, the accelerator compiler **104A-C** may compile and/or otherwise generate one(s) of the machine-learning model(s) **124** as lightweight machine-learning models.

[0062] In general, implementing an ML/AI system involves two phases, a learning/training phase and an inference phase. In the learning/training phase, a training algorithm is used to train the machine-learning model(s) **124** to operate in accordance with patterns and/or associations based on, for example, training data. In general, the machine-learning model(s) **124** include(s) internal parameters (e.g., the configuration data **122**) that guide how input data is transformed into output data, such as through a series of nodes and connections within the machine-learning model(s) **124** to transform input data into output data. Additionally, hyperparameters (e.g., the configuration data **122**) are used as part of the training process to control how the learning is performed (e.g., a learning rate, a number of layers to be used in the machine learning model, etc.). Hyperparameters are defined to be training parameters that are determined prior to initiating the training process.

[0063] Different types of training may be performed based on the type of ML/AI model and/or the expected output. For example, the accelerator compiler **104A-C** may invoke supervised training to use inputs and corresponding expected (e.g., labeled) outputs to select parameters (e.g., by iterating over combinations of select parameters) for the machine-learning model(s) **124** that reduce model error. As used herein, “labeling” refers to an expected output of the machine learning model (e.g., a classification, an expected output value, etc.). Alternatively, the accelerator compiler **104A-C** may invoke unsupervised training (e.g., used in deep learning, a subset of machine learning, etc.) that

involves inferring patterns from inputs to select parameters for the machine-learning model(s) **124** (e.g., without the benefit of expected (e.g., labeled) outputs).

[0064] In some examples, the accelerator compiler **104A-C** trains the machine-learning model(s) **124** using unsupervised clustering of operating observables. However, the accelerator compiler **104A-C** may additionally or alternatively use any other training algorithm such as stochastic gradient descent, Simulated Annealing, Particle Swarm Optimization, Evolution Algorithms, Genetic Algorithms, Nonlinear Conjugate Gradient, etc.

[0065] In some examples, the accelerator compiler **104A-C** may train the machine-learning model(s) **124** until the level of error is no longer reducing. In some examples, the accelerator compiler **104A-C** may train the machine-learning model(s) **124** locally on the computing system **102** and/or remotely at an external computing system (e.g., one(s) of the external computing systems **130**) communicatively coupled to the computing system **102**. In some examples, the accelerator compiler **104A-C** trains the machine-learning model(s) **124** using hyperparameters that control how the learning is performed (e.g., a learning rate, a number of layers to be used in the machine learning model, etc.). In some examples, the accelerator compiler **104A-C** may use hyperparameters that control model performance and training speed such as the learning rate and regularization parameter(s). The accelerator compiler **104A-C** may select such hyperparameters by, for example, trial and error to reach an optimal model performance. In some examples, the accelerator compiler **104A-C** utilizes Bayesian hyperparameter optimization to determine an optimal and/or otherwise improved or more efficient network architecture to avoid model overfitting and improve the overall applicability of the machine-learning model(s) **124**.

Alternatively, the accelerator compiler **104A-C** may use any other type of optimization. In some examples, the accelerator compiler **104A-C** may perform re-training. The accelerator compiler **104A-C** may execute such re-training in response to override(s) by a user of the computing system **102**, a receipt of new training data, etc.

[0066] In some examples, the accelerator compiler **104A-C** facilitates the training of the machine-learning model(s) **124** using training data. In some examples, the accelerator compiler **104A-C** utilizes training data that originates from locally generated data. In some examples, the accelerator compiler **104A-C** utilizes training data that originates from externally generated data. In some examples where supervised training is used, the accelerator compiler **104A-C** may label the training data. Labeling is applied to the training data by a user manually or by an automated data pre-processing system. In some examples, the accelerator compiler **104A-C** may pre-process the training data using, for example, an interface (e.g., the interface circuitry **114**). In some examples, the accelerator compiler **104A-C** sub-divides the training data into a first portion of data for training the machine-learning model(s) **124**, and a second portion of data for validating the machine-learning model(s) **124**.

[0067] Once training is complete, the accelerator compiler **104A-C** may deploy the machine-learning model(s) **124** for use as an executable construct that processes an input and provides an output based on the network of nodes and connections defined in the machine-learning model(s) **124**. The accelerator compiler **104A-C** may store the machine-learning model(s) **124** in the datastore **120**. In some examples, the accelerator compiler **104A-C** may invoke the interface circuitry **114** to transmit the machine-learning model(s) **124** to one(s) of the external computing systems **130**. In some such examples, in response to transmitting the machine-learning model(s) **124** to the one(s) of the external computing systems **130**, the one(s) of the external computing systems **130** may execute the machine-learning model(s) **124** to execute AI/ML workloads with at least one of improved efficiency or performance.

[0068] Once trained, the deployed one(s) of the machine-learning model(s) **124** may be operated in an inference phase to process data. In the inference phase, data to be analyzed (e.g., live data) is input to the machine-learning model(s) **124**, and the machine-learning model(s) **124** execute(s) to create an output. This inference phase can be thought of as the AI “thinking” to generate the output

based on what it learned from the training (e.g., by executing the machine-learning model(s) **124** to apply the learned patterns and/or associations to the live data). In some examples, input data undergoes pre-processing before being used as an input to the machine-learning model(s) **124**. Moreover, in some examples, the output data may undergo post-processing after it is generated by the machine-learning model(s) **124** to transform the output into a useful result (e.g., a display of data, a detection and/or identification of an object, an instruction to be executed by a machine, etc.).

[0069] In some examples, output of the deployed one(s) of the machine-learning model(s) **124** may be captured and provided as feedback. By analyzing the feedback, an accuracy of the deployed one(s) of the machine-learning model(s) **124** can be determined. If the feedback indicates that the accuracy of the deployed model is less than a threshold or other criterion, training of an updated model can be triggered using the feedback and an updated training data set, hyperparameters, etc., to generate an updated, deployed model.

[0070] In some examples, the accelerator compiler **104A-C** configures one(s) of the acceleration circuitry **108, 110** to execute a convolution operation, such as 2-D convolution operation. For example, the acceleration circuitry **108, 110** may implement a CNN. In some examples, CNNs ingest and/or otherwise process images as tensors, which are matrices of numbers with additional dimensions. For example, a CNN can obtain an input image represented by 3-D tensors, where a first and a second dimension correspond to a width and a height of a matrix and a third dimension corresponds to a depth of the matrix. For example, the width and the height of the matrix can correspond to a width and a height of an input image and the depth of the matrix can correspond to a color depth (e.g., a color layer) or a color encoding of the image (e.g., a Red-Green-Blue (RGB) encoding).

[0071] A typical CNN may also receive an input and transform the input through a series of hidden layers. For example, a CNN may have a plurality of convolution layers, pooling layers, and/or fully-connected layers. In some such examples, a CNN may have a plurality of layer triplets including a convolution layer, a pooling layer, and a fully-connected layer. In some examples, a CNN may have a plurality of convolution and pooling layer pairs that output to one or more fully-connected layers. In some examples, a CNN may include 20 layers, 30 layers, etc.

[0072] In some examples, the acceleration circuitry **108, 110** may execute a convolution layer to apply a convolution function or operation to map images of an input (previous) layer to the next layer in a CNN. In some examples, the convolution may be three-dimensional (3-D) because each input layer can have multiple input features (e.g., input channels) associated with an input image. The acceleration circuitry **108, 110** may execute the convolution layer to perform convolution by forming a regional filter window in each individual input channel and generating output data or activations by calculating a product of (1) a filter weight associated with the regional filter window and (2) the input data covered by the regional filter window. For example, the acceleration circuitry **108, 110** may determine an output feature of an input image by using the convolution filter to scan a plurality of input channels including a plurality of the regional filter windows.

[0073] In some examples, the acceleration circuitry **108, 110** may execute a pooling layer to extract information from a set of activations in each output channel. The pooling layer may perform a maximum pooling operation corresponding to a maximum pooling layer or an average pooling operation corresponding to an average pooling layer. In some examples, the maximum pooling operation may include selecting a maximum value of activations within a pooling window. In some examples, the average pooling operation may include calculating an average value of the activations within the pooling window.

[0074] In some examples, the acceleration circuitry **108, 110** may execute a fully-connected layer to obtain the data calculated by the convolution layer(s) and/or the pooling layer(s) and/or classify the data into one or more classes. In some examples, the fully-connected layer may determine whether the classified data corresponds to a particular image feature of the input image. For

example, the acceleration circuitry **108, 110** may execute the fully-connected layer to determine whether the classified data corresponds to a simple image feature (e.g., a horizontal line) or a more complex image feature like an animal (e.g., a cat).

[0075] In some examples, the accelerator compiler **104A-C** may configure one(s) of the acceleration circuitry **108, 110** to execute non-2-D convolution operations as 2-D convolution operations. For example, the accelerator compiler **104A-C** may configure the one(s) of the acceleration circuitry **108, 110** to implement a depthwise convolution operation, an elementwise addition operation, a grouped convolution operation, a dilated convolution operation, a custom operation (e.g., a custom convolution, a custom acceleration operation, etc.), etc., as a 2-D convolution operation. In some such examples, the accelerator compiler **104A-C** may instruct the one(s) of the acceleration circuitry **108, 110** to internally generate data rather than receive the data from the accelerator compiler **104A-C**, the configuration data **122**, etc. For example, the accelerator compiler **104A-C** may instruct the first acceleration resource to generate at least one of activation sparsity data, weight sparsity data, or weight data based on the acceleration operation to be executed by the first acceleration circuitry **108**. In some such examples, the accelerator compiler **104A-C** may instruct the one(s) of the acceleration circuitry **108, 110** to execute the one(s) of the ML model(s) **124** based on the data generated by the one(s) of the acceleration circuitry **108, 110**, which may be based on a convolution operation to be executed by the one(s) of the acceleration circuitry **108, 110**.

[0076] FIG. 2 is an illustration of an example accelerator **200**. In some examples, the accelerator **200** may implement the first acceleration circuitry **108** and/or the second acceleration circuitry **110** of FIG. 1. For example, the accelerator **200** may implement an AI and/or ML accelerator (e.g., a GPU, a VPU, an NN processor, etc.). The accelerator **200** of FIG. 2 is a hardware accelerator. For example, the accelerator **200** may be implemented by one or more analog or digital circuit(s), logic circuits, programmable processor(s), programmable controller(s), GPU(s), DSP(s), ASIC(s), PLD(s), and/or FPLD(s)). In some examples, the accelerator **200** implements a machine-learning model, such as a NN, or portion(s) thereof. For example, the accelerator **200** may implement one of the ML model(s) **124** of FIG. 1, or one or more layers (e.g., ML layers, NN layers, etc.) thereof.

[0077] In the illustrated example of FIG. 2, the accelerator **200** is a sparse accelerator, such as a sparse NN accelerator. For example, an accelerator that implements a CNN may execute a convolution layer using one or more multiply-accumulate (MAC) operations between input activations and kernel weights. The accelerator **200** may save energy and/or otherwise execute with reduced power consumption (and reduce the time taken to process a NN layer) by ignoring and/or otherwise not calculating activation values having a value of zero. In some examples, the accelerator **200** may implement a ReLU function as the non-linear activation. For example, the ReLU function may clamp all negative values to zero, which may result in a significant number of zeros being present in the output activations and, thus, may be the main source of activation sparsity.

[0078] In some examples, the accelerator **200** may implement weight sparsity in response to a pruning of small magnitude values and replacing them with a value of zero. For example, the accelerator compiler **104A-C** of FIG. 1 may introduce sparsity by employing techniques such as certain types of regularization to encourage weight values to zero during training of the machine-learning model (e.g., the ML model(s) **124** of FIG. 1) to be implemented by the accelerator **200**.

[0079] As described in further detail below, the accelerator **200** may combine both activation and weight sparsity to fully leverage all redundancy to achieve improved and/or otherwise maximum acceleration.

[0080] An accelerator based on a dense NNA may only read in data, such as activation and weight data. A sparse accelerator based on a sparse NNA may read in both the data and control information. In some examples, the control information may be used to indicate where the non-zero data is located in the read-in data. In some examples, the control information may be implemented

in different control formats to represent the sparse data, such as run-length-encoded bit streams, coordinate lists, bit masks of non-zero entries, etc. In the example of FIG. 2, the accelerator **200** may use a bit mask of non-zero entries (e.g., a sparsity map or a sparsity vector) to represent the sparse data.

[0081] The accelerator **200** of the illustrated example includes accelerator or acceleration inputs, which may include an example configuration interface **202**, an example activation data interface (I/F) **204**, an example activation sparsity data interface **206**, an example weight data interface **208**, an example weight sparsity data interface **210**. The accelerator **200** includes accelerator or acceleration outputs, which may include an example output activation interface **212**. For example, at least one of the configuration interface **202**, the activation data interface **204**, the activation sparsity data interface **206**, the weight data interface **208**, the weight sparsity data interface **210**, or the output activation interface **212** may be implemented by one or more interface circuits (e.g., interface circuitry, interface logic circuitry, etc.) implemented by any type of interface standard. In some such examples, at least one of the configuration interface **202**, the activation data interface **204**, the activation sparsity data interface **206**, the weight data interface **208**, the weight sparsity data interface **210**, or the output activation interface **212** may be implemented by one or more analog or digital circuit(s), logic circuits, etc.

[0082] The accelerator **200** of the illustrated example includes example configuration circuitry **214**, an example activation data buffer **216**, an example external activation sparsity buffer **218**, example internal activation sparsity generation circuitry **220**, an example weight data buffer **222**, example internal weight generation circuitry **224**, example external weight sparsity buffer **226**, example internal weight sparsity generation circuitry **228**, a first example multiplexer (MUX) **230**, a second example multiplexer **232**, an example combined sparsity controller **234**, a plurality of example multiply accumulators (MACs) **236**, and example accumulator storage **238**. In some examples, other structures may be present in the accelerator **200**. For example, the accelerator **200** may include direct memory access (DMA) circuitry, CPU scheduling circuitry, etc., or any other type of accelerator circuitry. In some such examples, the accelerator **200** is not so limited to the structures depicted in the illustrated example of FIG. 2.

[0083] Output(s) of the activation data interface **204** is/are coupled to input(s) (e.g., buffer input(s), data buffer input(s), etc.) of the activation data buffer **216**. Output(s) (e.g., interface output(s)) of the configuration interface **202** is/are coupled to input(s) (e.g., configuration input(s), configuration circuitry input(s), etc.) of the configuration circuitry **214**. Output(s) (e.g., configuration output(s), configuration circuitry output(s), etc.) of the configuration circuitry **214** is/are coupled to input(s) of the internal activation sparsity generation circuitry **220**, input(s) of the internal weight sparsity generation circuitry **228**, input(s) (e.g., selection input(s), control input(s), multiplexer control input(s), etc.) of the first multiplexer **230**, and input(s) (e.g., selection input(s), control input(s), multiplexer control input(s), etc.) of the second multiplexer **232**.

[0084] In this example, output(s) of the activation sparsity data interface **206** is/are coupled to input(s) of the external activation sparsity buffer **218**. Output(s) of the internal activation sparsity generation circuitry **220** and output(s) (e.g., buffer output(s), data buffer output(s), etc.) of the external activation sparsity buffer **218** is/are coupled to respective input(s) (e.g., multiplexer input(s)) of the first multiplexer **230**. Output(s) (e.g., multiplexer output(s)) of the first multiplexer **230** is/are coupled to input(s) (e.g., controller input(s)) of the combined sparsity controller **234**.

[0085] In this example, output(s) of the weight data interface **208** is/are coupled to input(s) of the weight data buffer **222**. Output(s) of the weight sparsity data interface **210** is/are coupled to input(s) of the external weight sparsity buffer **226**. Output(s) of the internal weight sparsity generation circuitry **228** and output(s) of the external weight sparsity buffer **226** is/are coupled to respective input(s) of the second multiplexer **232**. Output(s) of the second multiplexer **232** is/are coupled to input(s) of the combined sparsity controller **234**. Output(s) (e.g., controller output(s)) of the combined sparsity controller **234** is/are coupled to input(s) of the activation data buffer **216** and

input(s) of the weight data buffer **222**.

[0086] In this example, output(s) of the activation data buffer **216** is/are coupled to input(s) of the MACs **236**. For example, one or more first outputs of the activation data buffer **216** may be coupled to one or more first inputs of a first one of the MACs **236**, one or more second outputs of the activation data buffer **216** may be coupled to one or more second inputs of a second one of the MACs **236**, etc. For example, the output(s) of the activation data buffer **216** may be coupled to respective ones of the accelerator inputs, which may be implemented by inputs of the MACs **236**.

[0087] In this example, output(s) of the weight data buffer **222** is/are coupled to input(s) of the MACs **236**. For example, one or more first outputs of the weight data buffer **222** may be coupled to one or more first inputs of a first one of the MACs **236**, one or more second outputs of the weight data buffer **222** may be coupled to one or more second inputs of a second one of the MACs **236**, etc. Output(s) of the MACs **236** is/are coupled to input(s) (e.g., storage input(s), accumulator storage input(s), etc.) of the accumulator storage **238**. Output(s) (e.g., storage output(s), accumulator storage output(s), etc.) of the accumulator storage **238** is/are coupled to input(s) of the output activation interface **212**. Output(s) of the output activation interface **212** may be coupled to different hardware. For example, the output(s) of the output activation interface **212** may be coupled to one(s) of the CPU **106**, the first acceleration circuitry **108**, the second acceleration circuitry **110**, the general purpose processing circuitry **112**, the interface circuitry **114**, the bus **116**, the datastore **120**, etc., of FIG. **1**.

[0088] In the illustrated example of FIG. **2**, the accelerator **200** includes the configuration interface **202**, the activation data interface **204**, the activation sparsity data interface **206**, the weight data interface **208**, and the weight sparsity data interface **210** to receive and/or otherwise obtain data from a data source (e.g., hardware, software, and/or firmware of the computing system **102** or portion(s) thereof). In some examples, the configuration interface **202** receives configuration data, such as the configuration data **122** of FIG. **1**. For example, the configuration interface **202** may receive the configuration data **122** to include configuration register settings. In some such examples, the configuration register settings may include a first configuration register setting that, when invoked, instructs the internal activation sparsity generation circuitry **220** to generate an activation bit mask (e.g., an activation sparsity bit mask) based on a first pattern (e.g., a first data pattern). In some such examples, the configuration register settings may include a second configuration register setting that, when invoked, instructs the internal weight sparsity generation circuitry **228** to generate a weight bit mask (e.g., a weight sparsity bit mask) based on a second pattern (e.g., a second data pattern).

[0089] In some examples, the activation data interface **204** receives activation data (e.g., convolution activation values, ML activation values, etc.) and stores the activation data in the activation data buffer **216**. For example, the external activation sparsity buffer **218** may be implemented by volatile memory (e.g., flash memory), non-volatile memory, etc. In some examples, the activation sparsity data interface **206** receives sparsity activation data (e.g., a sparsity activation vector) and stores the sparsity activation data in the external activation sparsity buffer **218**. In some examples, the weight data interface **208** receives weight data (e.g., convolution weight values, ML weight values, etc.) and stores the weight data in the weight data buffer **222**. In some examples, the weight sparsity data interface **210** receives sparsity weight data (e.g., a sparsity weight vector) and stores the sparsity weight data in the external weight sparsity buffer **226**. For example, the external weight sparsity buffer **226** may be implemented by volatile memory, non-volatile memory, etc.

[0090] In the illustrated example of FIG. **2**, the accelerator **200** includes the internal activation sparsity generation circuitry **220** to generate activation sparsity data based on configuration data, such as the configuration data **122** of FIG. **1**. For example, the internal activation sparsity generation circuitry **220** may generate a sparse activation vector based on a data pattern, which may be indicated and/or otherwise based on the configuration data.

[0091] In the illustrated example of FIG. 2, the accelerator **200** includes the internal weight sparsity generation circuitry **228** to generate weight sparsity data based on configuration data, such as the configuration data **122** of FIG. 1. For example, the internal weight sparsity generation circuitry **228** may generate a sparse weight vector based on a data pattern, which may be indicated and/or otherwise based on the configuration data.

[0092] Advantageously, the internal activation sparsity generation circuitry **220** and/or the internal weight sparsity generation circuitry **228** increases and/or otherwise improves the bandwidth of the accelerator **200**. For example, in a conventional accelerator architecture, sparsity vectors reside in memory of the accelerator. For a conventional convolution operation, the corresponding bit masks may be computed offline for weights or at runtime for activations. The bit masks may not be calculated internally within the conventional accelerator architecture for a conventional convolution because the distribution of zeros within the data is random.

[0093] Advantageously, for other types of accelerator operations in which the bit mask vectors may be indicative of compute (e.g., a '1' to indicate computation of an element and a '0' to indicate a non-computation of an element) as opposed to sparsity (e.g., a '1' to indicate a non-zero value and a '0' to indicate a zero value), the internal activation sparsity generation circuitry **220** and/or the internal weight sparsity generation circuitry **228** may calculate and/or otherwise generate bit masks to follow a fixed or repeating pattern (e.g., fixed or repeating data pattern). For example, the internal activation sparsity generation circuitry **220** and/or the internal weight sparsity generation circuitry **228** may generate the bit masks (e.g., the activation sparsity bit mask, the weight sparsity bit mask, etc.) internally and/or otherwise within the accelerator **200** rather than retrieving the bit masks from hardware external to the accelerator **200**. In some such examples, the internal activation sparsity generation circuitry **220** and/or the internal weight sparsity generation circuitry **228** may include driving hardware logic to produce, generate, and/or otherwise output data patterns, such as walking ones (1s) or repeating patterns of x is followed by y 0s. For example, a walking is pattern may be implemented by a first bit mask vector having a size of 1 byte of "00000001," a second bit mask vector having a size of 1 byte of "00000010," etc., where the "1" of each subsequent vector is left-shifted into an adjacent bit position of the previous bit position.

[0094] In some examples, the bit mask vectors (e.g., the activation sparsity bit mask vectors, the weight sparsity bit mask vectors, etc.) may use 1 bit of control to represent 1 byte (or 8 bits) of data. In some such examples, there is a minimum overhead of 12.5% (e.g., $0.125 = 1 \text{ bit of control} / 8 \text{ bits of data}$) for the accelerator **200** to read in the control information along with the data.

Advantageously, by generating the bit mask vectors internally with the internal activation sparsity generation circuitry **220** and/or the internal weight sparsity generation circuitry **228**, the accelerator **200** may remove the above-referenced overhead. In some examples, the removal of the overhead is beneficial when processing bandwidth limited layers such as depthwise convolution as the reduction of the overhead may reduce the number of reads to process the layer.

[0095] In the illustrated example of FIG. 2, the accelerator **200** includes the first multiplexer **230** to multiplex, select, and/or otherwise provide activation sparsity data from either the external activation sparsity buffer **218** or the internal activation sparsity generation circuitry **220**. For example, the configuration circuitry **214** may instruct the first multiplexer **230** to select first activation sparsity data from the external activation sparsity buffer **218** based on the configuration data, which may indicate a type of convolution operation to be executed by the accelerator **200**. In some such examples, the configuration circuitry **214** may instruct the first multiplexer **230** to select second activation sparsity data from the internal activation sparsity generation circuitry **220** based on the configuration data, which may indicate a type of convolution operation to be executed by the accelerator **200**. In response to the instruction from the configuration circuitry **214**, the first multiplexer **230** may provide the first activation sparsity data or the second activation sparsity data to the combined sparsity controller **234**.

[0096] In the illustrated example of FIG. 2, the accelerator **200** includes the second multiplexer **232**

to multiplex, select, and/or otherwise provide weight sparsity data from either the external weight sparsity buffer **226** or the internal weight sparsity generation circuitry **228**. For example, the configuration circuitry **214** may instruct the second multiplexer **232** to select first weight sparsity data from the external weight sparsity buffer **226** based on the configuration data, which may indicate a type of convolution operation to be executed by the accelerator **200**. In some such examples, the configuration circuitry **214** may instruct the second multiplexer **232** to select second weight sparsity data from the internal weight sparsity generation circuitry **228** based on the configuration data, which may indicate a type of convolution operation to be executed by the accelerator **200**. For example, the configuration circuitry **214** may determine an accelerator mode based on the type of acceleration operation to be executed by the accelerator **200**. In some such examples, the configuration circuitry **214** may determine the acceleration mode to be a 2-D convolution mode, a depthwise convolution mode, an elementwise addition mode, a grouped convolution mode, a dilated convolution mode, a custom mode (e.g., a mode in which the accelerator **200** is to execute a custom operation), etc., based on the configuration data. In response to the instruction from the configuration circuitry **214**, the second multiplexer **232** may provide the first weight sparsity data or the second weight sparsity data to the combined sparsity controller **234**. In some examples, the first multiplexer **230** and/or the second multiplexer **232** may implement multiplexer circuitry, multiplexer logic circuitry, etc.

[0097] In the illustrated example of FIG. 2, the accelerator **200** includes the combined sparsity controller **234** to provide and/or otherwise transfer non-sparse data from at least one of the activation data buffer **216** or the weight data buffer **222** to the MACs **236**. For example, the combined sparsity controller **234** may implement control circuitry, control logic circuitry, controller circuitry, controller logic circuitry, etc. In some examples, the combined sparsity controller **234** determines that the presence of zeros in activation or weight data is indicative of when to skip certain computations when performing an acceleration operation, such as a convolution. For example, the combined sparsity controller **234** may determine that a sparsity bit mask vector is indicative of a “compute” bit mask vector where each bit may denote whether an element should be computed or non-computed (e.g., skipped). In some such examples, the combined sparsity controller **234** may determine that a ‘1’ in a bit mask indicates that an element should be part of a computation operation and a ‘0’ in a bit mask indicates that an element should not be part of a computation operation. In some such examples, the computation operation may be a standard or conventional convolution operation (e.g., a 2-D convolution operation), a non-standard or non-conventional convolution operation (e.g., a depthwise convolution operation, a grouped convolution operation, etc.), or a non-convolution operation (e.g., an elementwise addition operation).

[0098] In some examples, the combined sparsity controller **234** identifies the non-sparse data to be transferred based on a sparsity map. For example, the combined sparsity controller **234** may identify non-sparse activation data based on the sparse activation vector obtained from the first multiplexer **230**. In some such examples, the combined sparsity controller **234** may map a ‘1’ in the sparse activation vector to a non-sparse activation value stored in the activation data buffer **216**. In some such examples, the activation data buffer **216** transfers the identified non-sparse activation value to one of the MACs **236**.

[0099] In some examples, the combined sparsity controller **234** may identify non-sparse weight data based on the sparse weight vector obtained from the second multiplexer **232**. In some such examples, the combined sparsity controller **234** may map a ‘1’ in the sparse weight vector to a non-sparse weight value stored in the weight data buffer **222**. In some such examples, the weight data buffer **222** transfers the identified non-sparse weight value to one of the MACs **236**. In some examples, the weight data buffer **222** may invoke the internal weight generation circuitry **224** to generate the non-sparse weight value, which may be transferred to the one of the MACs **236**.

[0100] In some examples, the combined sparsity controller **234** obtains sparsity data from at least

one of the first multiplexer **230** or the second multiplexer **232**. For example, the sparsity data may include activation sparsity data from the first multiplexer **230** and/or weight sparsity data from the second multiplexer **232**. In some examples, the combined sparsity controller **234** may generate a combined sparsity bit mask based on the activation sparsity data and the weight sparsity data. In some such examples, the combined sparsity controller **234** may instruct the activation data buffer **216** to provide activation data to the MACs **236** based on the combined sparsity bit mask. In some examples, the combined sparsity controller **234** may instruct the weight data buffer **222** to provide weight data to the MACs **236** based on the combined sparsity bit mask. In some such examples, in response to a determination that the internal weight generation circuitry **224** is to generate weight data, the combined sparsity controller **234** may instruct the weight data buffer **222** to provide the internally generated weight data to the MACs **236** based on the combined sparsity bit mask.

[0101] In the illustrated example of FIG. 2, the accelerator **200** includes the MACs **236** to execute and/or otherwise implement a multiply-accumulate operation. For example, a first one of the MACs **236** may determine the product of two numbers and add that product to an accumulator. In some examples, the MACs **236** may implement accelerator or acceleration logic (e.g., accelerator or acceleration circuitry, accelerator or acceleration logic circuitry, etc.), compute or computation logic (e.g., compute or computation circuitry, compute or computation logic circuitry, etc.), etc. In this example, the MACs **236** are implemented as a 4×4 array of the MACs **236**. Alternatively, the accelerator **200** may implement the MACs **236** in a different manner. For example, the accelerator **200** may implement the MACs **236** as one or more 8×8 arrays, one or more 16×16 arrays, etc., or any other quantity and/or type of array configuration(s). In this example, the MACs **236** may follow a common interconnection where activations may be shared and distributed across rows of the array and weights may be shared and distributed along columns of the array. In some examples, each of the MACs **236** may have an associated accumulator storage that may allow partial and/or final accumulations to be saved and thereby enable the MACs **236** to work on multiple output points per workload. In some such examples, in response to completing a workload (e.g., a convolution workload, an acceleration or accelerator workload, etc.), the final accumulated values may be drained and/or otherwise output from the accumulator storage **238** to external hardware via the output activation interface **212** for final post-processing. In some examples, the final post-processing may include adding a bias and/or applying a non-linear activation function before writing the post-processed values to memory (e.g., an external memory). For example, the accumulator storage **238** may provide the final accumulated values to the memory via the output activation interface **212**. In some examples, the accumulator storage **238** may output the final accumulated values as corresponding to a layer of a NN. In some such examples, in response to completing the workload for the layer, which may be indicated by the output from the accumulator storage **238**, the accelerator **200** may process another layer of the NN until the entire NN is complete.

[0102] Advantageously, the accelerator **200** of FIG. 2 may utilize sparsity (e.g., activation sparsity and/or weight sparsity) for faster processing and/or reduced power consumption. For example, the accelerator **200** may combine activation and weight sparsity to remove and/or skip redundant computation to achieve faster and/or otherwise improved processing of layers of a machine-learning model (e.g., a machine-learning network), reduce power consumption, and/or provide sparse acceleration. Advantageously, the accelerator **200** may reduce the cost of data movement (e.g., a bandwidth penalty) and/or reduce the bandwidth requirement for reading in activations and/or weights by removing activations and/or weights that have a value of zero to improve memory storage density (e.g., increased storage of non-zero data values compared to zero data values). Advantageously, in some examples, the accelerator **200** may increase and/or otherwise improve bandwidth by internally generating data (e.g., activation sparsity data, weight data, weight sparsity data, etc.) rather than retrieving the data from an external source.

[0103] In some examples, the internal activation sparsity generation circuitry **220** and/or the

internal weight sparsity generation circuitry **228** may implement example means for generating sparsity data based on an acceleration operation. For example, the means for generating may be implemented by executable instructions such as that implemented by at least blocks **1404**, **1408**, **1412** of FIG. **14**, blocks **1510**, **1512**, **1516** of FIG. **15**, blocks **1608**, **1612**, **1616**, **1620**, and/or block **1714** of FIG. **17**. In some examples, the executable instructions of blocks **1404**, **1408**, **1412** of FIG. **14**, blocks **1510**, **1512**, **1516** of FIG. **15**, blocks **1608**, **1612**, **1616**, **1620**, and/or block **1714** of FIG. **17** may be executed on at least one processor such as the example processor circuitry **1812** of FIG. **18**. In other examples, the means for generating is implemented by hardware logic, hardware implemented state machines, logic circuitry, and/or any other combination of hardware, software, and/or firmware. For example, the means for generating may be implemented by at least one hardware circuit (e.g., discrete and/or integrated analog and/or digital circuitry, an FPGA, a PLD, a FPLD, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware, but other structures are likewise appropriate.

[0104] In some examples in which sparsity data includes a sparse weight vector, the means for generating is to generate a sparsity bit mask based on a sparse activation vector and the sparse weight vector in response to obtaining the sparse activation vector. In some examples in which the sparsity data includes a sparse weight vector, and the acceleration operation is a depthwise convolution operation, the means for generating is to identify a data pattern based on the depthwise convolution operation, and generate the sparse weight vector based on the data pattern.

[0105] In some examples in which the sparsity data includes a sparse weight vector, and the acceleration operation is a grouped convolution operation, the means for generating is to identify a data pattern based on the grouped convolution operation, the data pattern including a set of logic ones, and generate the sparse weight vector based on the data pattern.

[0106] In some examples in which the acceleration operation is an elementwise addition operation, the means for generating is to identify a data pattern based on the elementwise addition operation, the data pattern based on a walking ones data pattern, and generate a plurality of sparse weight vectors based on the data pattern.

[0107] In some examples in which the sparsity data includes one or more sparse weight vectors, and the acceleration operation is a dilated convolution, the means for generating is to identify a data pattern based on the dilated convolution, the data pattern based on a sparse kernel, and generate the one or more sparse weight vectors based on the data pattern.

[0108] In some examples, the combined sparsity controller **234** may implement example means for instructing one or more data buffers to provide at least one of activation data or weight data based on sparsity data. For example, the means for instructing may be implemented by executable instructions such as that implemented by at least blocks **1412**, **1416** of FIG. **14**, blocks **1516**, **1518** of FIG. **15**, and/or block **1622** of FIG. **16**. In some examples, the executable instructions of blocks **1412**, **1416** of FIG. **14**, blocks **1516**, **1518** of FIG. **15**, and/or block **1622** of FIG. **16** may be executed on at least one processor such as the example processor circuitry **1812** of FIG. **18**. In other examples, the means for instructing is implemented by hardware logic, hardware implemented state machines, logic circuitry, and/or any other combination of hardware, software, and/or firmware. For example, the means for instructing may be implemented by at least one hardware circuit (e.g., discrete and/or integrated analog and/or digital circuitry, an FPGA, a PLD, a FPLD, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware, but other structures are likewise appropriate.

[0109] In some examples in which the sparsity data includes at least one of activation sparsity data or weight sparsity data, and the one or more data buffers include a weight data buffer and an activation data buffer, the means for instructing is to generate a combined sparsity bit mask based on the activation sparsity data and the weight sparsity data, instruct the activation data buffer to

provide activation data to the means for executing based on the combined sparsity bit mask, and, in response to determining to generate weight data, instruct the weight data buffer to provide the generated weight data to the means for executing based on the combined sparsity bit mask.

[0110] In some examples in which the sparsity data includes a sparse weight vector, and the acceleration operation is a depthwise convolution operation, the means for instructing is to generate a sparsity bit mask based on the sparse weight vector. In some examples in which the sparsity data includes a sparse weight vector, and the acceleration operation is a grouped convolution operation, the means for instructing is to generate a sparsity bit mask based on the sparse weight vector. In some examples in which the sparsity data includes one or more sparse weight vectors, and the acceleration operation is a dilated convolution operation, the means for instructing is to generate one or more sparsity bit masks based on the one or more sparse weight vectors.

[0111] In some examples, the MACs **236** may implement example means for executing an acceleration operation based on at least one of activation data or weight data. For example, the means for instructing may be implemented by executable instructions such as that implemented by at least block **1416** of FIG. **15**, block **1518** of FIG. **15**, and/or block **1622** of FIG. **16**. In some examples, the executable instructions of block **1416** of FIG. **15**, block **1518** of FIG. **15**, and/or block **1622** of FIG. **16** may be executed on at least one processor such as the example processor circuitry **1812** of FIG. **18**. In other examples, the means for executing is implemented by hardware logic, hardware implemented state machines, logic circuitry, and/or any other combination of hardware, software, and/or firmware. For example, the means for executing may be implemented by at least one hardware circuit (e.g., discrete and/or integrated analog and/or digital circuitry, an FPGA, a PLD, a FPLD, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware, but other structures are likewise appropriate.

[0112] In some examples in which the acceleration operation is an elementwise addition operation, the means for executing is to store first data values associated with a first tensor based on the plurality of the sparse weight vectors and a first set of the activation data, store second data values associated with a second tensor based on the plurality of the sparse weight vectors and a second set of the activation data, and generate output data values based on a sum of the first data values and the second data values.

[0113] In some examples in which sparsity data includes a sparse weight vector, the configuration circuitry **214** may implement example means for identifying an acceleration operation based on configuration information, and the sparse weight vector may be generated based on the configuration information. For example, the means for identifying may be implemented by executable instructions such as that implemented by at least blocks **1402**, **1404**, **1408**, **1412**, **1420** of FIG. **14**, blocks **1502**, **1504**, **1510**, **1516**, **1522** of FIG. **15**, and/or blocks **1602**, **1606**, **1608**, **1610**, **1612**, **1614**, **1616**, **1618**, **1620**, **1624** of FIG. **16**. In some examples, the executable instructions of blocks **1402**, **1404**, **1408**, **1412**, **1420** of FIG. **14**, blocks **1502**, **1504**, **1510**, **1516**, **1522** of FIG. **15**, and/or blocks **1602**, **1606**, **1608**, **1610**, **1612**, **1614**, **1616**, **1618**, **1620**, **1624** of FIG. **16** may be executed on at least one processor such as the example processor circuitry **1812** of FIG. **18**. In other examples, the means for identifying is implemented by hardware logic, hardware implemented state machines, logic circuitry, and/or any other combination of hardware, software, and/or firmware. For example, the means for identifying may be implemented by at least one hardware circuit (e.g., discrete and/or integrated analog and/or digital circuitry, an FPGA, a PLD, a FPLD, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware, but other structures are likewise appropriate.

[0114] While an example manner of implementing the first example acceleration circuitry **108** and/or the second example acceleration circuitry **110** of FIG. **1** is illustrated in FIG. **2**, one or more of the elements, processes and/or devices illustrated in FIG. **2** may be combined, divided, re-

arranged, omitted, eliminated and/or implemented in any other way. Further, the example configuration interface **202**, the example activation data interface **204**, the example activation sparsity data interface **206**, the example weight data interface **208**, the example weight sparsity data interface **210**, the example output activation interface **212**, the example configuration circuitry **214**, the example activation data buffer **216**, the example external activation sparsity buffer **218**, the example internal activation sparsity generation circuitry **220**, the example weight data buffer **222**, the example internal weight generation circuitry **224**, the example external weight sparsity buffer **226**, the example internal weight sparsity generation circuitry **228**, the first example multiplexer **230**, the second example multiplexer **232**, the example combined sparsity controller **234**, the example MACs **236**, the example accumulator storage **238**, and/or, more generally, the first example acceleration circuitry **108** and/or the second example acceleration circuitry **110** of FIG. **1** may be implemented by hardware, software, firmware, and/or any combination of hardware, software and/or firmware. Thus, for example, any of the example configuration interface **202**, the example activation data interface **204**, the example activation sparsity data interface **206**, the example weight data interface **208**, the example weight sparsity data interface **210**, the example output activation interface **212**, the example configuration circuitry **214**, the example activation data buffer **216**, the example external activation sparsity buffer **218**, the example internal activation sparsity generation circuitry **220**, the example weight data buffer **222**, the example internal weight generation circuitry **224**, the example external weight sparsity buffer **226**, the example internal weight sparsity generation circuitry **228**, the first example multiplexer **230**, the second example multiplexer **232**, the example combined sparsity controller **234**, the example MACs **236**, the example accumulator storage **238**, and/or, more generally, the first example acceleration circuitry **108** and/or the second example acceleration circuitry **110** could be implemented by one or more analog or digital circuit(s), logic circuits, programmable processor(s), programmable controller(s), GPU(s), DSP(s), ASIC(s), PLD(s), and/or FPLD(s). When reading any of the apparatus or system claims of this patent to cover a purely software and/or firmware implementation, at least one of the example configuration interface **202**, the example activation data interface **204**, the example activation sparsity data interface **206**, the example weight data interface **208**, the example weight sparsity data interface **210**, the example output activation interface **212**, the example configuration circuitry **214**, the example activation data buffer **216**, the example external activation sparsity buffer **218**, the example internal activation sparsity generation circuitry **220**, the example weight data buffer **222**, the example internal weight generation circuitry **224**, the example external weight sparsity buffer **226**, the example internal weight sparsity generation circuitry **228**, the first example multiplexer **230**, the second example multiplexer **232**, the example combined sparsity controller **234**, the example MACs **236**, and/or the example accumulator storage **238** is/are hereby expressly defined to include a non-transitory computer readable storage device or storage disk such as a memory, a DVD, a CD, a Blu-ray disk, etc., including the software and/or firmware. Further still, the first example acceleration circuitry **108** and/or the second acceleration circuitry **110** of FIG. **1** may include one or more elements, processes and/or devices in addition to, or instead of, those illustrated in FIG. **2**, and/or may include more than one of any or all of the illustrated elements, processes and devices. As used herein, the phrase “in communication,” including variations thereof, encompasses direct communication and/or indirect communication through one or more intermediary components, and does not require direct physical (e.g., wired) communication and/or constant communication, but rather additionally includes selective communication at periodic intervals, scheduled intervals, aperiodic intervals, and/or one-time events.

[0115] FIG. **3** is a block diagram of an example accelerator compiler **300**. In some examples, the accelerator compiler **300** of FIG. **3** may implement one or more of the accelerator compiler **104A-C** of FIG. **1**. The accelerator compiler **300** may configure a hardware accelerator, such as the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2** based on an accelerator mode. For example, the accelerator compiler **300**

may configure the accelerator **200** of FIG. 2 based on the accelerator **200** to execute a 2-D convolution operation, an element-wise addition operation, etc.

[0116] The accelerator compiler **300** of the illustrated example includes an example machine-learning (ML) model configuration handler **310**, an example activation configuration handler **320**, an example weight configuration handler **330**, an example configuration compiler **340**, an example executable generator **350**, an example accelerator interface **360**, an example machine-learning (ML) model datastore **370**, an example activation datastore **380**, an example weight datastore **390**, and an example bus **395**. In this example, the ML model configuration handler **310**, the activation configuration handler **320**, the weight configuration handler **330**, the configuration compiler **340**, the executable generator **350**, the accelerator interface **360**, the ML model datastore **370**, the activation datastore **380**, and the weight datastore **390** are in communication with one (s0 of each other via the bus **395**. For example, the bus **395** may be implemented by at least one of an I2C bus, a SPI bus, a PCI bus, or a PCIe bus. Additionally or alternatively, the bus **395** may be implemented by any other type of computing or electrical bus.

[0117] In the illustrated example of FIG. 3, the accelerator compiler **300** includes the ML model configuration handler **310** to configure an accelerator, such as the first acceleration circuitry **108** of FIG. 1, the second accelerator circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2, based on an acceleration operation to be executed by the accelerator, and/or an acceleration mode of the accelerator. In some examples, the ML model configuration handler **310** determines the acceleration mode based on the acceleration operation to be executed by the accelerator. For example, the ML model configuration handler **310** may determine that the accelerator is to execute a depthwise convolution operation based on a type of NN to be executed by the accelerator. In some such examples, the type of NN may indicate that the depthwise convolution operation is to be executed. In some such examples, the ML model configuration handler **310** may determine the acceleration mode to be a depthwise convolution mode based on the depthwise convolution operation to be executed by the accelerator.

[0118] In the illustrated example of FIG. 3, the accelerator compiler **300** includes the activation configuration handler **320** to identify activation data and/or activation sparsity data to be utilized by the accelerator. In some examples, the activation configuration handler **320** may identify activation data to be provided to the activation data buffer **216** and/or activation sparsity data to be provided to the external activation sparsity buffer **218**. In some examples, the activation configuration handler **320** may determine that the accelerator **200** is to generate the activation sparsity data based on the acceleration mode. In some such examples, the activation configuration handler **320** may identify a value of one or more configuration registers (e.g., configuration control registers, etc.) that may be utilized to instruct the internal activation sparsity generation circuitry **220** to generate activation sparsity data based on a data pattern.

[0119] In the illustrated example of FIG. 3, the accelerator compiler **300** includes the weight configuration handler **330** to identify weight data and/or weight sparsity data to be utilized by the accelerator. In some examples, the weight configuration handler **330** may identify weight data to be provided to the weight data buffer **222** and/or weight sparsity data to be provided to the external weight sparsity buffer **226**.

[0120] In some examples, the weight configuration handler **330** may determine that the accelerator **200** is to generate the weight data and/or the weight sparsity data based on the acceleration mode. In some such examples, the weight configuration handler **330** may identify a value of one or more configuration registers (e.g., configuration control registers, etc.) that may be utilized to instruct the internal weight sparsity generation circuitry **228** to generate weight sparsity data based on a first data pattern, the internal weight generation circuitry **224** to generate weight data based on a second data pattern, etc., and/or a combination thereof.

[0121] In the illustrated example of FIG. 3, the accelerator compiler **300** includes the configuration compiler **340** to compile an ML, network, etc., based on configuration data. In some examples, the

configuration compiler **340** may compile the ML model, the network, etc., based on the configuration data **122** of FIG. **1**, which may include values of one or more configuration registers, which may be utilized by the accelerator **200** to internally generate weight data, weight sparsity data, and/or activation sparsity data within the accelerator **200**.

[0122] In the illustrated example of FIG. **3**, the accelerator compiler **300** includes the executable generator **350** to generate an executable construct of the ML model, the network, etc., based on the compilation of the configurations by the configuration compiler **340**. In some examples, the executable generator **350** may generate the executable construct to be executed by the accelerator to output accumulation values from the accumulator storage **238** based on ML inputs to at least one of the activation data interface **204**, the activation sparsity data interface **206**, the weight data interface **208**, or the weight sparsity data interface **210**. In some examples, the executable generator **350** may generate the executable construct as the ML model(s) **124** of FIG. **1**.

[0123] In the illustrated example of FIG. **3**, the accelerator compiler **300** includes the accelerator interface **360** to interface with an accelerator. In some examples, the accelerator interface **360** may provide the configuration data **122** to the configuration interface **202**. In some examples, the accelerator interface **360** may provide activation data to the activation data interface **204**. In some examples, the accelerator interface **360** may provide weight data to the weight data interface **208**. In some examples, the accelerator interface **360** may provide weight sparsity data to the weight sparsity data interface **210**. In some examples, the accelerator interface **360** may obtain accumulator data (e.g., accumulator values from the MACs **236** of FIG. **2**) from the output activation interface **212**.

[0124] In the illustrated example of FIG. **3**, the accelerator compiler **300** includes the ML model datastore **370** to store one or more ML models, such as the ML model(s) **124** of FIG. **1**. In some examples, the ML model datastore **370** may implement the ML model(s) **124**, and/or, more generally, the datastore **120** of FIG. **1**. In the illustrated example of FIG. **3**, the accelerator compiler **300** includes the activation datastore **380** to record data, such as activation data, activation sparsity data, etc. In the illustrated example of FIG. **3**, the accelerator compiler **300** includes the weight datastore **390** to record data, such as weight data, weight activation data, etc.

[0125] The ML model datastore **370**, the activation datastore **380**, and/or the weight datastore **390** of this example may be implemented by a volatile memory and/or a non-volatile memory. The ML model datastore **370**, the activation datastore **380**, and/or the weight datastore **390** may additionally or alternatively be implemented by one or more DDR memories. The ML model datastore **370**, the activation datastore **380**, and/or the weight datastore **390** may additionally or alternatively be implemented by one or more mass storage devices. While in the illustrated example the ML model datastore **370**, the activation datastore **380**, and the weight datastore **390** are each illustrated as single datastores, the ML model datastore **370**, the activation datastore **380**, and/or the weight datastore **390** may be implemented by any number and/or type(s) of datastores. Furthermore, the data stored in the ML model datastore **370**, the activation datastore **380**, and/or the weight datastore **390** may be in any data format such as, for example, binary data, comma delimited data, tab delimited data, SQL structures, an executable, etc.

[0126] While an example manner of implementing the accelerator compiler **104A-C** of FIG. **1** is illustrated in FIG. **3**, one or more of the elements, processes and/or devices illustrated in FIG. **3** may be combined, divided, re-arranged, omitted, eliminated and/or implemented in any other way. Further, the example ML model configuration handler **310**, the example activation configuration handler **320**, the example weight configuration handler **330**, the example configuration compiler **340**, the example executable generator **350**, the example accelerator interface **360**, the example ML model datastore **370**, the example activation datastore **380**, the example weight datastore **390**, the example bus **395**, and/or, more generally, the example accelerator compiler **104A-C** of FIG. **1** may be implemented by hardware, software, firmware and/or any combination of hardware, software and/or firmware. Thus, for example, any of the example ML model configuration handler **310**, the

example activation configuration handler **320**, the example weight configuration handler **330**, the example configuration compiler **340**, the example executable generator **350**, the example accelerator interface **360**, the example ML model datastore **370**, the example activation datastore **380**, the example weight datastore **390**, the example bus **395**, and/or, more generally, the example accelerator compiler **104A-C** of FIG. **1** could be implemented by one or more analog or digital circuit(s), logic circuits, programmable processor(s), programmable controller(s), GPU(s), DSP(s), ASIC(s), PLD(s), and/or FPLD(s). When reading any of the apparatus or system claims of this patent to cover a purely software and/or firmware implementation, at least one of the example ML model configuration handler **310**, the example activation configuration handler **320**, the example weight configuration handler **330**, the example configuration compiler **340**, the example executable generator **350**, the example accelerator interface **360**, the example ML model datastore **370**, the example activation datastore **380**, the example weight datastore **390**, and/or the example bus **395** is/are hereby expressly defined to include a non-transitory computer readable storage device or storage disk such as a memory, a DVD, a CD, a Blu-ray disk, etc., including the software and/or firmware. Further still, the example accelerator compiler **104A-C** of FIG. **1** may include one or more elements, processes and/or devices in addition to, or instead of, those illustrated in FIG. **2**, and/or may include more than one of any or all of the illustrated elements, processes and devices.

[0127] FIG. **4** is an illustration of an example conventional convolution operation **400** that may be executed by the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2**. In some examples, the conventional convolution operation **400** may implement a spatial convolution over one or more images (e.g., a picture, a still frame of a video, etc.). In some examples, the accelerator **200** may be configured to operate in a conventional convolution mode, a 2-D convolution mode, a three-dimensional (3-D) convolution mode, etc., based on the conventional convolution operation to be executed by the accelerator **200**.

[0128] The conventional convolution operation **400** includes applying example filters **402** to an example input tensor **404** to generate an example output tensor **406**. In this example, the input tensor **404** is a 3-D object having a size of $x_i * y_i * z_i$. In this example, there are K of the filters **402** and each of the filters **402** have a size of $f * f * z_{sub.k}$. Alternatively, any other size may be used to implement one(s) of the filters **402**. For example, one or more of the filters **402** may have a size of $f_{sub.x} * f_{sub.y} * z_{sub.k}$ where x and y may be different and thereby $f_{sub.x}$ and $f_{sub.y}$ may be different. In this example, the filters **402** are square filters and thereby $f_{sub.x}$ is equal to $f_{sub.y}$ but examples described herein are not so limited. In this example, the output tensor **406** has a size of $x_o * y_o * z_o$. In this example, $z_k = z_i$ and $z_o = K$. In this example, the filters **402** along with a non-linear activation function are applied to the input tensor **404** to produce the output tensor **406**. For example, the accelerator **200** of FIG. **2** may obtain the input tensor **404** as the activation data, obtain one of the filters **402** as the weight data, and output the output tensor **406** from the accumulator storage **238** of FIG. **2**. In some such examples, the accelerator **200** may implement the conventional convolution operation **400** in a “dense” manner while, in other examples, the accelerator **200** may implement the conventional convolution operation **400** utilizing sparsity.

[0129] Advantageously, the accelerator **200** may execute the convolution operation **400** based on sparse data to reduce the number of computations. For example, the accelerator **200** may obtain and/or generate activation sparsity data and/or weight sparsity data to output the output tensor **406** by invoking sparsity techniques.

[0130] FIG. **5** is an illustration of an example depthwise convolution operation **500** that may be executed by the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2**. The depthwise convolution operation **500** includes applying example filters **502** to an example input tensor **504** to generate an example output tensor **506**.

[0131] The illustrated example of FIG. **5** is an example of utilizing bit mask vectors (e.g., sparsity bit mask vectors) to isolate a single channel at a time. For example, the configuration circuitry **214**

may determine to invoke the accelerator **200** to execute the depthwise convolution operation **500** of FIG. **5** based on the configuration data **122** of FIG. **1**. In some such examples, the configuration circuitry **214** may instruct the internal weight sparsity generation circuitry **228** to generate the weight sparsity bit masks for the depthwise convolution operation **500** based on the configuration register settings, which may be based on the configuration data **122**. In some such examples, the internal weight sparsity generation circuitry **228** may generate the weight sparsity bit masks based on a data pattern, such as a walking is pattern. For example, the internal weight sparsity generation circuitry **228** may generate the filters **502** based on a walking is pattern that, when utilized, causes the filters **502** to generate a result (e.g., the output tensor **506**) that may be equivalent to a depthwise filter, which may allow the accelerator **200** to execute sparse conventional convolutions to implement depthwise convolution.

[0132] The depthwise convolution operation **500** may be implemented as a spatial convolution performed independently over every example input channel **508** within the input tensor **504**. The depthwise convolution operation **500** of the illustrated example may be based on a determination that a depthwise convolution may be executed on a sparse neural network accelerator, such as the accelerator **200**, by configuring the accelerator **200** in the depthwise convolution mode. For example, by considering the filters **502** as sparse filters with only a single non-sparse bit, the depthwise convolution operation **500** may be implemented using the accelerator **200**, which is depicted in the illustrated example of FIG. **6**.

[0133] FIG. **6** is an illustration of example sparsity bit mask vectors **602**, **604**, **606**, **608** associated with a convolution filter. In this example, the sparsity bit mask vectors **602**, **604**, **606**, **608** may be implemented by the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2**. The sparsity bit mask vectors **602**, **604**, **606**, **608** include first example sparsity bit mask vectors **602**, second example sparsity bit mask vectors **604**, third example sparsity bit mask vectors **606**, and fourth example sparsity bit mask vectors **608**. For example, the first sparsity bit mask vectors **602**, the second sparsity bit mask vectors **604**, the third sparsity bit mask vectors **606**, or the fourth sparsity bit mask vectors **608** may implement the filters **402** of FIG. **4** and/or the filters **502** of FIG. **5**. In some such examples, the internal weight sparsity generation circuitry **228** of FIG. **2** may generate the first sparsity bit mask vectors **602**, the second sparsity bit mask vectors **604**, the third sparsity bit mask vectors **606**, and/or the fourth sparsity bit mask vectors **608**. In some such examples, the internal weight sparsity generation circuitry **228** of FIG. **2** may generate the first sparsity bit mask vectors **602**, the second sparsity bit mask vectors **604**, the third sparsity bit mask vectors **606**, and/or the fourth sparsity bit mask vectors **608** based on the configuration data **122**, a data pattern, etc.

[0134] The first sparsity bit mask vectors **602** may implement fully dense sparsity vectors that have a filter length of zK . The second sparsity bit mask vectors **604** may implement sparse sparsity vectors, which may include a combination of 1s and 0s due to the presence of the sparsity. In this example, the second sparsity bit mask vectors **604** have an effective filter length less than zK .

[0135] The third sparsity bit mask vectors **606** of the illustrated example may implement walking is sparsity vectors that have an effective filter length of 1. For example, the third sparsity bit mask vectors **606** may be generated based on a data pattern, such as the walking ones data pattern (e.g., walking is data pattern). The fourth sparsity bit mask vectors **608** may implement sparsity bit mask vectors associated with depthwise filters that have an effective filter depth zK of 1. In some such examples, the walking is data pattern of third sparsity bit mask vectors **606** may be considered a subset of the sparse vector example depicted as the second sparsity bit mask vectors **604** and thereby may be utilized to mimic and/or otherwise be operative as a depthwise kernel or filter. For example, the third sparsity bit mask vectors **606** may be equivalent and/or otherwise substantially approximate to the fourth sparsity bit mask vectors **608** because they have the same effective length. Because depthwise convolutions may be bandwidth intensive, the internal weight sparsity generation circuitry **228** may generate the third sparsity bit mask vectors **606** internal to the

accelerator **200** to eliminate the need to read in the sparsity patterns for the third sparsity bit mask vectors **606** to thereby save bandwidth and reduce convolution time.

[0136] Advantageously, by instructing the internal weight sparsity generation circuitry **228** to generate the third sparsity bit mask vectors **606**, the accelerator **200** may execute the depthwise convolution operation **500** of FIG. 5 without dedicated depthwise convolution hardware. By reducing and/or otherwise eliminating convolution-specific hardware, such as specific depthwise convolution hardware, the accelerator **200** may implement a plurality of convolution operation types, which may significantly ease physical implementation of the accelerator **200**, reduce area of the accelerator **200**, and/or enhance scalability of convolution operations to be executed by one or more of the accelerator **200**.

[0137] FIG. 7 is an illustration of an example grouped convolution operation **700** that may be executed by the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2. The illustrated example of FIG. 7 is an example of utilizing bit mask vectors (e.g., sparsity bit mask vectors) to isolate a number of channels at a time and then perform the grouped convolution operation **700** as a conventional convolution operation. In some examples, grouped convolutions may be implemented by performing a number of convolutions across groups of channels within a layer instead of a single convolution across all channels within the layer.

[0138] Further depicted in the illustrated example of FIG. 7 is a non-grouped convolution operation **702**, which processes a first example input layer **704** of a NN by applying 64 example filters **706** to the 256 channels of the first input layer to generate a first example output layer **708**. The layers of the grouped convolution operation **700** and the non-grouped convolution operation **702** are depicted with a format of (#input channels, filter size, #output channels). For example, the input layer **704** has 256 input channels, a filter size of 1×1 , and 64 output channels.

[0139] The grouped convolution operation **700** of the illustrated example of FIG. 7 may be implemented by applying groups of 4 example filters **710** to a second example input layer **712** to create groups of example 4-layer outputs **714**. The 4-layer outputs **714** may be added together output a second example output layer **716**. In the grouped convolution operation **700**, there are 32 groups of 4-layer deep tensors in the middle row of the example. The shallow and grouped processing may include executing more smaller workloads as opposed to single bigger workloads depicted by the non-grouped convolution operation **702**.

[0140] In some examples, bigger, deeper workloads may be more efficient and/or otherwise more optimal for NNAs to process as opposed to smaller, shallower workloads. Advantageously, the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2 may utilize manipulation of bit masks to cause the accelerator **200** to execute the second input layer **712** (e.g., the 32 (256, 1×1 , 4) convolutions) as a single convolution (e.g., (256, 1×1 , 128) convolution)) and yield a 128-d output. In some examples, the accelerator **200** may process the 128-d output as a single tensor (e.g., (128, 3×3 , 128) tensor) and yield the same output as the 32 (4, 3×3 , 4) convolutions of the grouped convolution operation **700**.

[0141] FIG. 8 is an illustration of example bit masks **802**, **804**, **806**, **808**, **810** that may be utilized to implement the grouped convolution operation **700** of FIG. 7, which may be executed by the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2. In this example, the bit masks **802**, **804**, **806**, **808**, **810** may implement sparsity bit masks (e.g., activation sparsity bit masks, weight sparsity bit masks, etc.).

[0142] In some examples, the bit masks **802**, **804**, **806**, **808**, **810** may be generated by the internal activation sparsity generation circuitry **220** and/or the internal weight sparsity generation circuitry **228** of FIG. 2 based on the configuration data **122** of FIG. 1. For example, the configuration circuitry **214** may determine to invoke the accelerator **200** to execute the grouped convolution operation **700** of FIG. 7 based on the configuration data **122** of FIG. 1. In some such examples, the configuration circuitry **214** may instruct the internal weight sparsity generation circuitry **228** to

generate the bit masks **802, 804, 806, 808, 810** for the grouped convolution operation **700** based on the configuration register settings, which may be based on the configuration data **122**. In some such examples, the internal weight sparsity generation circuitry **228** may generate the bit masks **802, 804, 806, 808, 810** based on an example data pattern **812**. In this example, the data pattern **812** is a grouped walking 1s pattern, which includes a set of ones (e.g., a set of logic ones), a group of ones (e.g., a group of logic ones), etc.

[0143] The bit masks **802, 804, 806, 808, 810** include 32 bit masks having a filter length of $z=256$, which include a first example bit mask **802** corresponding to Group 0, a second example bit mask **804** corresponding to Group 1, a third example bit mask **806** corresponding to Group 2, a fourth example bit mask **808** corresponding to Group 3, and a fifth example bit mask **810** corresponding to Group 31. For example, the 32 ones of the bit masks **802, 804, 806, 808, 810** may be utilized to isolate 4 channels at a time from the 256-channel deep tensor of FIG. 7 (e.g., 256-d in). In some such examples, the 0s in the bit masks **802, 804, 806, 808, 810** may ensure that the filter only effectively sees 4 channels. Advantageously, the accelerator **200** may manipulate the bit masks **802, 804, 806, 808, 810** based on the data pattern **812** (or a different data pattern) to execute the grouped convolution operation **700** with increased efficiency

[0144] FIG. 9 is an illustration of a first example elementwise addition operation **900** that may be executed by the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2. The first elementwise addition operation **900** may be an example of utilizing bit mask manipulation and a structure of a MAC array to isolate a single activation at a time.

[0145] The first elementwise addition operation **900** includes an example input tensor **902**, example rows of activation buffers **904**, example columns of weight buffers **906**, an example MAC array **908**, an example shared accumulator storage **910** for the MAC array, and example output activations **912**. The input tensor **902** is a $4 \times 1 \times 4$ tensor. Alternatively, the input tensor **902** may be any other size. The rows of activation buffers **904** may include activation data stored in the activation data buffer **216** of FIG. 2. The columns of weight buffers **906** may include weight data stored in the weight data buffer **222** of FIG. 2. For example, the internal weight sparsity generation circuitry **228** may generate the weight data of the columns of weight buffers **906** based on the configuration data **122** of FIG. 1. The MAC array **908** may be an example implementation of ones of the MACs **236** of FIG. 2. In this example, the MAC array **236** is a 4×4 MAC array.

Alternatively, the MAC array **236** may be any other size and/or configuration. In this example, the MACs **236** are provided weight data. For example, the internal weight generation circuitry **224** may generate the weight data based on the configuration data **122**. The shared accumulator storage **910** may implement the accumulator storage **238** of FIG. 2. The output activations **912** may include activation data stored in the accumulator storage **238** of FIG. 2.

[0146] Advantageously, an accelerator described herein, such as the accelerator **200**, may utilize internal generation of sparsity patterns and weight values plus internal storage to execute the first elementwise addition operation **900** using the MAC array **908**, which may typically be utilized for conventional convolution. Advantageously, the accelerator **200** may reconfigure hardware tailored for conventional convolution to implement non-convolution operations, such as the first elementwise addition operation **900**, which may be performed with negligible area cost.

[0147] In some examples, the accelerator **200** may determine that a sparsity vector can be used to isolate a single element in a tensor, such as the input tensor **902**, for elementwise processing. For example, in addition to isolating a single activation, the weight data may be internally generated by the accelerator **200** and hardcoded to a value (e.g., 0x1) that may be fixed in hardware and/or configured by firmware and/or software. In some such examples, the isolated activation value may be stored in the shared accumulator storage **910** when the isolated activation value passes through the MAC array **908**, and/or, more generally, the accelerator **200**.

[0148] In the illustrated example of FIG. 9, the MAC array **908** may be utilized to obtain a $4 \times 1 \times 4$

input tensor as an input isolate the 16×1 activations based on the walking is bit mask and the operation of the MAC array **908**. In this example, the MAC array **908** have a configuration where rows of the MAC array **908** share activation buffers and columns of the MAC array **908** share weight buffers. In the top row of this example, all of the MACs in the first row of the MAC array **908** see $y[0]x[0]z[3:0]$ of the input tensor **902**, which, when the bit mask vectors of the columns of weight buffers **906** are applied and the 0×1 multiplication is performed, the 4 accumulator values in the top row are populated with $y[0]x[0]z[0]$, $y[0]x[0]z[1]$, $y[0]x[0]z[2]$, and $y[0]x[0]z[3]$. Advantageously, the generation of the bit masks based on a data pattern of a walking is pattern may effectuate the isolation of the single activation values within the input tensor **902**.

[0149] FIG. **10** is an illustration of a second example elementwise addition operation **1000** that may be executed by the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2**. In this example, the second elementwise addition operation **1000** may be implemented by the addition of a first example tensor (TENSOR A) **1002** and a second example tensor (TENSOR B) **1004** to generate and/or otherwise output a third example tensor (TENSOR C) **1006**.

[0150] In this example, the first tensor **1002** of FIG. **10** may implement the input tensor **902** of FIG. **9** or portion(s) thereof. For example, the first elementwise addition operation **900** of FIG. **9** may be executed to isolate first activation values of the first tensor **1002** in a first acceleration operation and isolate second activation values of the second tensor **1004** in a second acceleration operation. In some such examples, the first elementwise addition operation **900** of FIG. **9** may be executed to add the first activation values and the second activation values to output third activation values of the third tensor **1006**.

[0151] In some examples, the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2** may execute (e.g., iteratively execute) the first elementwise addition operation **900** of FIG. **9** to achieve an addition of the first tensor **1002** and the second tensor **1004** of FIG. **10**. In some such examples, the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2** may execute (e.g., iteratively execute) the first elementwise addition operation **900** of FIG. **9** by breaking down the first tensor **1002** and the second tensor **1004** into smaller blocks or tensors and add (e.g., iteratively add) together to generate the output as depicted in the illustrated example of FIG. **11**.

[0152] FIG. **11** is an illustration of a third example elementwise addition operation **1100** that may be executed by the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2**. In this example, the first tensor **1002** and the second tensor **1004** may be divided, partitioned, and/or otherwise broken up into smaller tensors or blocks. In this example, the first tensor **1002** and the second tensor **1004** may be divided into blocks of $16 \times 8 \times 16$. In this example, a first block of $16 \times 8 \times 16$ of the first tensor **1002** may be added to a second block of $16 \times 8 \times 16$ of the second tensor **1004** to yield a third $16 \times 8 \times 16$ block of the third tensor **1006**. For example, the accelerator **200** of FIG. **2** may execute the first elementwise addition operation **900** of FIG. **9** on the first block and the second block to yield the third block of the third tensor **1006**. In some such examples, the accelerator **200** may repeat and/or otherwise iteratively execute the first elementwise addition operation **900** 16 times on the first tensor **1002** and the second tensor **1004** to output the $64 \times 16 \times 32$ block of the third tensor **1006**.

[0153] Advantageously, the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2** may perform the first elementwise addition operation **900**, the second elementwise addition operation **1000**, and/or the third elementwise addition operation **1100**. For example, the accelerator **200** may be implemented with hardware to execute conventional convolution operations. In some such examples, the accelerator **200** may be configured to execute the conventional convolution operation, the first elementwise addition operation **900**, the second elementwise addition operation **1000**, and/or the third

elementwise addition operation **1100** with the same hardware, which may reduce the hardware cost of supporting and/or otherwise implementing elementwise addition operations.

[0154] Advantageously, the first acceleration circuitry **108**, the second acceleration circuitry **110**, and/or the accelerator **200** may execute the first elementwise addition operation **900**, the second elementwise addition operation **1000**, and/or the third elementwise addition operation **1100** with dense and/or sparse vectors. For example, the accelerator **200** may execute the first elementwise addition operation **900**, the second elementwise addition operation **1000**, and/or the third elementwise addition operation **1100** to effectuate the addition of two sparse vectors, a sparse vector and a dense vector, or a two dense vectors utilizing the same hardware of FIG. 2.

[0155] FIG. 12 is an illustration of example dilated convolution operations **1202**, **1204**, **1206** that may be executed by the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2. The dilated convolution operations **1202**, **1204**, **1206** of FIG. 12 may implement examples of generating bit mask vectors to manipulate a structure of filter kernels to perform the dilated convolution operations **1202**, **1204**, **1206** as conventional 2-D convolution operations. As depicted in the illustrated example, dilated convolutions, or atrous convolutions, may expand a receptive field of a first example kernel **1208** by inserting holes between adjacent elements.

[0156] In this example, the dilated convolution is a 3×3 dilated convolution. The illustrated example depicts three example dilated convolutions **1202**, **1204**, **1206** including a first example dilated convolution **1202**, a second example dilated convolution **1204**, and a third example dilated convolution **1206**. The first dilated convolution **1202** may be implemented with a dilation of 1 in which there are no holes between adjacent elements of a second example kernel **1210**. The second dilated convolution **1204** may be implemented with a dilation of 2 in which there is a single hole between adjacent elements of a third example kernel **1212**. The third dilated convolution **1206** may be implemented with a dilation of 3 in which there are two holes between adjacent elements of a fourth example kernel **1214**.

[0157] FIG. 13 is an illustration of example kernels **1302**, **1304** that may be implemented by the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, and/or the accelerator **200** of FIG. 2. For example, the internal activation sparsity generation circuitry **220** of FIG. 2 may generate the kernels **1302**, **1304**. In some examples, the internal weight sparsity generation circuitry **228** of FIG. 2 may generate the kernels **1302**, **1304**.

[0158] In the illustrated example of FIG. 13, the kernels **1302**, **1304** include a first example kernel **1302** and a second example kernel **1304**. In this example, the first kernel **1302** is a 5×5 dense kernel and the second kernel **1304** is a 5×5 sparse kernel. For example, the internal weight sparsity generation circuitry **228** may generate the second kernel **1304** to implement a 3×3 dilated convolution operation based on a dilation of 2. In some such examples, the first kernel **1302** may implement the second kernel **1210** of FIG. 12 and/or the second kernel **1304** may implement the third kernel **1212** of FIG. 12. For example, the second kernel **1304** may be considered a sparse subset of a large, non-dilated kernel, such as the first kernel **1302**. Advantageously, the internal weight sparsity generation circuitry **228** may manipulate, modify, and/or otherwise adjust the first kernel **1302** by inserting 0s between adjacent elements to translate and/or otherwise convert the first kernel **1302** into the second kernel **1304**. In some examples, the internal weight sparsity generation circuitry **228** may generate the second kernel **1304** by programming configuration registers of the configuration circuitry **214** and/or the internal weight sparsity generation circuitry **228** to produce the repeating pattern internally or by manipulating the values in memory of the configuration circuitry **214** and/or the internal weight sparsity generation circuitry **228** and allowing the weight data buffer **222** to read such values. Advantageously, the internal weight sparsity generation circuitry **228**, and/or, more generally, the accelerator **200** may effectuate a dilated convolution operation, such as one of the dilated convolution operations **1202**, **1204**, **1206** of FIG. 12, to be performed as a conventional convolution operation on a sparse NNA, such as the

accelerator **200**, by manipulating the first kernel **1302** to generate the second kernel **1304**. [0159] Flowcharts representative of example hardware logic, machine readable instructions, hardware implemented state machines, and/or any combination thereof for implementing the example accelerator compiler **104A-C** of FIG. **1**, the first example acceleration circuitry **108** of FIG. **1**, the second example acceleration circuitry **110** of FIG. **1**, the example accelerator **200** of FIG. **2**, and/or the accelerator compiler **300** of FIG. **3** are shown in FIGS. **14-17**. The machine readable instructions may be one or more executable programs or portion(s) of an executable program for execution by a computer processor and/or processor circuitry, such as the example processor circuitry **1812**, the example graphics processing unit **1840**, the example vision processing unit **1842**, and/or the example neural network processor **1844** shown in the example processor platform **1800** discussed below in connection with FIG. **18**, the processor circuitry discussed below in connection with FIG. **19**, and/or the processor circuitry discussed below in connection with FIG. **20**. The program may be embodied in software stored on a non-transitory computer readable storage medium such as a CD-ROM, a floppy disk, an HDD, a DVD, a Blu-ray disk, or a memory associated with the example processor circuitry **1812**, the example graphics processing unit **1840**, the example vision processing unit **1842**, and/or the example neural network processor **1844**, but the entire program and/or parts thereof could alternatively be executed by a device other than the example processor circuitry **1812**, the example graphics processing unit **1840**, the example vision processing unit **1842**, the example neural network processor **1844**, and/or embodied in firmware or dedicated hardware. Further, although the example program is described with reference to the flowcharts illustrated in FIGS. **14-17**, many other methods of implementing the example accelerator compiler **104A-C** of FIG. **1**, the first example acceleration circuitry **108** of FIG. **1**, the second example acceleration circuitry **110** of FIG. **1**, the example accelerator **200** of FIG. **2**, and/or the accelerator compiler **300** of FIG. **3** may alternatively be used. For example, the order of execution of the blocks may be changed, and/or some of the blocks described may be changed, eliminated, or combined. Additionally or alternatively, any or all of the blocks may be implemented by one or more hardware circuits (e.g., discrete and/or integrated analog and/or digital circuitry, an FPGA, an ASIC, a comparator, an operational-amplifier (op-amp), a logic circuit, etc.) structured to perform the corresponding operation without executing software or firmware. The processor circuitry may be distributed in different network locations and/or local to one or more devices (e.g., a multi-core processor in a single machine, multiple processors distributed across a server rack, etc.).

[0160] The machine readable instructions described herein may be stored in one or more of a compressed format, an encrypted format, a fragmented format, a compiled format, an executable format, a packaged format, etc. Machine readable instructions as described herein may be stored as data or a data structure (e.g., portions of instructions, code, representations of code, etc.) that may be utilized to create, manufacture, and/or produce machine executable instructions. For example, the machine readable instructions may be fragmented and stored on one or more storage devices and/or computing devices (e.g., servers) located at the same or different locations of a network or collection of networks (e.g., in the cloud, in edge devices, etc.). The machine readable instructions may require one or more of installation, modification, adaptation, updating, combining, supplementing, configuring, decryption, decompression, unpacking, distribution, reassignment, compilation, etc. in order to make them directly readable, interpretable, and/or executable by a computing device and/or other machine. For example, the machine readable instructions may be stored in multiple parts, which are individually compressed, encrypted, and stored on separate computing devices, wherein the parts when decrypted, decompressed, and combined form a set of executable instructions that implement one or more functions that may together form a program such as that described herein.

[0161] In another example, the machine readable instructions may be stored in a state in which they may be read by processor circuitry, but require addition of a library (e.g., a dynamic link library

(DLL)), a software development kit (SDK), an application programming interface (API), etc., in order to execute the instructions on a particular computing device or other device. In another example, the machine readable instructions may need to be configured (e.g., settings stored, data input, network addresses recorded, etc.) before the machine readable instructions and/or the corresponding program(s) can be executed in whole or in part. Thus, machine readable media, as used herein, may include machine readable instructions and/or program(s) regardless of the particular format or state of the machine readable instructions and/or program(s) when stored or otherwise at rest or in transit.

[0162] The machine readable instructions described herein can be represented by any past, present, or future instruction language, scripting language, programming language, etc. For example, the machine readable instructions may be represented using any of the following languages: C, C++, Java, C#, Perl, Python, JavaScript, HyperText Markup Language (HTML), Structured Query Language (SQL), Swift, etc.

[0163] As mentioned above, the example processes of FIGS. 14-17 may be implemented using executable instructions (e.g., computer and/or machine readable instructions) stored on a non-transitory computer and/or machine readable medium such as an HDD, a flash memory, a read-only memory, a CD, a DVD, a cache, a random-access memory, and/or any other storage device or storage disk in which information is stored for any duration (e.g., for extended time periods, permanently, for brief instances, for temporarily buffering, and/or for caching of the information). As used herein, the term non-transitory computer readable medium is expressly defined to include any type of computer readable storage device and/or storage disk and to exclude propagating signals and to exclude transmission media.

[0164] “Including” and “comprising” (and all forms and tenses thereof) are used herein to be open ended terms. Thus, whenever a claim employs any form of “include” or “comprise” (e.g., comprises, includes, comprising, including, having, etc.) as a preamble or within a claim recitation of any kind, it is to be understood that additional elements, terms, etc., may be present without falling outside the scope of the corresponding claim or recitation. As used herein, when the phrase “at least” is used as the transition term in, for example, a preamble of a claim, it is open-ended in the same manner as the term “comprising” and “including” are open ended. The term “and/or” when used, for example, in a form such as A, B, and/or C refers to any combination or subset of A, B, C such as (1) A alone, (2) B alone, (3) C alone, (4) A with B, (5) A with C, (6) B with C, and (7) A with B and with C. As used herein in the context of describing structures, components, items, objects and/or things, the phrase “at least one of A and B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B. Similarly, as used herein in the context of describing structures, components, items, objects and/or things, the phrase “at least one of A or B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B. As used herein in the context of describing the performance or execution of processes, instructions, actions, activities and/or steps, the phrase “at least one of A and B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B. Similarly, as used herein in the context of describing the performance or execution of processes, instructions, actions, activities and/or steps, the phrase “at least one of A or B” is intended to refer to implementations including any of (1) at least one A, (2) at least one B, and (3) at least one A and at least one B.

[0165] As used herein, singular references (e.g., “a”, “an”, “first”, “second”, etc.) do not exclude a plurality. The term “a” or “an” entity, as used herein, refers to one or more of that entity. The terms “a” (or “an”), “one or more”, and “at least one” can be used interchangeably herein. Furthermore, although individually listed, a plurality of means, elements or method actions may be implemented by, e.g., a single unit or processor. Additionally, although individual features may be included in different examples or claims, these may possibly be combined, and the inclusion in different

examples or claims does not imply that a combination of features is not feasible and/or advantageous.

[0166] FIG. 14 is a flowchart representative of an example process 1400 that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the example acceleration circuitry 108, 110 of FIG. 1 and/or the sparse accelerator 200 of FIG. 2 to execute an acceleration operation. The example process 1400 begins at block 1402, at which the first acceleration circuitry 108 and/or the second acceleration circuitry 110 identifies an acceleration operation based on configuration information associated with an accelerator. For example, the configuration circuitry 214 (FIG. 2) may identify the depthwise convolution operation 500 of FIG. 5 to be executed by the accelerator 200. In some such examples, the configuration circuitry 214 may identify the depthwise convolution operation 500 based on the configuration data 122 of FIG. 1.

[0167] At block 1404, the first acceleration circuitry 108 and/or the second acceleration circuitry 110 may configure the accelerator based on the acceleration operation. For example, the configuration circuitry 214 may configure one or more configuration registers of the internal activation sparsity generation circuitry 220 (FIG. 2), the internal weight sparsity generation circuitry 228 (FIG. 2), and/or the internal weight generation circuitry 224 (FIG. 2) based on the acceleration operation. In some such examples, the configuration circuitry 214 may configure the one or more configuration registers to generate a bit mask (e.g., an activation sparsity bit mask, a weight activation bit mask, etc.) and/or weight value(s) based on a data pattern, such as a walking is data pattern, a grouped walking is data pattern, etc., and/or a combination thereof.

[0168] At block 1406, the first acceleration circuitry 108 and/or the second acceleration circuitry 110 may obtain at least one of activation data, activation sparsity data, or weight data. For example, the activation data buffer 216 (FIG. 2) may obtain activation data from a resource (e.g., hardware, software, and/or firmware) of the computing system 102 of FIG. 1 via the activation data interface 204 (FIG. 2). In some such examples, the external activation sparsity buffer 218 (FIG. 2) may obtain activation sparsity data from a resource of the computing system 102 of FIG. 1 via the activation sparsity data interface 206 (FIG. 2). In some such examples, the weight data buffer 222 (FIG. 2) may obtain weight data from a resource of the computing system 102 of FIG. 1 via the weight data interface 208 (FIG. 2).

[0169] At block 1408, the first acceleration circuitry 108 and/or the second acceleration circuitry 110 may determine whether the accelerator is to generate weight sparsity data. For example, the configuration circuitry 214 may instruct the internal weight sparsity generation circuitry 228 (FIG. 2) to generate weight sparsity data based on the convolution acceleration operation to be executed by the accelerator 200.

[0170] If, at block 1408, the first acceleration circuitry 108 and/or the second acceleration circuitry 110 determine that the accelerator is not to generate the weight sparsity data, then, at block 1410, the first acceleration circuitry 108 and/or the second acceleration circuitry 110 obtain the weight sparsity data. For example, the external weight sparsity buffer 226 (FIG. 2) may obtain weight sparsity data from a resource of the computing system 102 of FIG. 1. In response to obtaining the weight sparsity data at block 1410, control proceeds to block 1414 to generate a sparsity bit mask based on the activation sparsity data and the weight sparsity data.

[0171] If, at block 1408, the first acceleration circuitry 108 and/or the second acceleration circuitry 110 determine that the accelerator is to generate the weight sparsity data, control proceeds to block 1412 to generate a sparsity bit mask based on the activation sparsity data and the weight sparsity data. For example, the configuration circuitry 214 may instruct the first multiplexer 230 to provide a sparse activation vector to the combined sparsity controller 234 (FIG. 2) based on the acceleration operation. In some examples, the configuration circuitry 214 may instruct the second multiplexer 232 to provide a sparse weight vector to the combined sparsity controller 234 based on the acceleration operation. In some such examples, the configuration circuitry 214 may instruct the

second multiplexer **232** to select the sparse weight vector from the internal weight sparsity generation circuitry **228** in response to an identification of the acceleration operation as a depthwise convolution operation. In some such examples, the combined sparsity controller **234** may generate a sparsity bit mask, such as the third sparsity bit mask vectors **606** of FIG. **6**, one of the bit masks **802, 804, 806, 808, 810**, etc., of FIG. **8**, etc., based on the activation sparsity data and the weight sparsity data.

[0172] At block **1416**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may instruct one or more data buffers to provide at least one of the activation data or the weight data based on the sparsity data to accelerator logic of the accelerator. For example, the combined sparsity controller **234** may instruct the activation data buffer **216** to provide activation data based on the sparsity bit mask, which may be based on the activation sparsity data, to the MACs **236**. In some examples, the combined sparsity controller **234** may instruct the weight data buffer **222** to provide weight data based on the sparsity bit mask, which may be based on the weight sparsity data, to the MACs **236**.

[0173] At block **1418**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may execute the acceleration operation based on the at least one of the activation data or the weight data. For example, the accelerator **200** may execute the depthwise convolution operation **500** of FIG. **5**, the grouped convolution operation **700** of FIG. **7**, the first elementwise addition operation **900** of FIG. **9**, the second elementwise addition operation **1000** of FIG. **10**, the third elementwise addition operation **1100** of FIG. **11**, the first dilated convolution operation **1202** of FIG. **12**, the second dilated convolution operation **1204** of FIG. **12**, the third dilated convolution operation **1206** of FIG. **12**, etc.

[0174] At block **1420**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether to select another acceleration operation to execute with the accelerator. For example, the configuration circuitry **214** may determine whether to select another acceleration operation to execute with the accelerator **200**. If, at block **1420**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine to select another acceleration operation to execute with the accelerator, control returns to block **1402** to identify the acceleration operation based on configuration information associated with the accelerator, otherwise the example process **1400** of FIG. **14** concludes.

[0175] FIG. **15** is a flowchart representative of an example process **1500** that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2** to execute an acceleration operation. The example process **1500** of FIG. **15** begins at block **1502**, at which the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may identify an acceleration operation based on configuration information associated with an accelerator. For example, the configuration circuitry **214** (FIG. **2**) may identify based on the configuration data **122** of FIG. **1** that the accelerator is to execute the depthwise convolution operation **500** of FIG. **5**, the first elementwise addition operation **900** of FIG. **9**, etc.

[0176] At block **1504**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the accelerator is to generate weight data. For example, the configuration circuitry **214** may determine based on the configuration data **122** that the internal weight generation circuitry **224** is to generate weight data. In some such examples, the configuration circuitry **214** may determine that the internal weight generation circuitry **224** is to generate weight data based on the configuration data **122** indicating that the acceleration operation to be executed by the accelerator **200** is the first elementwise addition operation **900** of FIG. **9**.

[0177] If, at block **1504**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the accelerator is not to generate weight data, control proceeds to block **1508**, otherwise, at block **1506**, the first acceleration circuitry **108** and/or the second acceleration circuitry

110 may generate weight data based on the acceleration operation. For example, the internal weight generation circuitry **224** may generate the weight data based on the acceleration operation being the first elementwise addition operation **900**.

[0178] At block **1508**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may receive at least one of activation or the weight data. For example, in response to a determination to generate the weight data, the activation data buffer **216** (FIG. 2) may receive activation data from a resource of the computing system **102** via the activation data interface **204** (FIG. 2). Advantageously, in some such examples, the generated weight data may not need to be received via the weight data interface **208** (FIG. 2). In some examples, in response to a determination not to generate the weight data, the activation data buffer **216** may receive activation data from a resource of the computing system **102** via the activation data interface **204** and the weight data buffer **222** may receive weight data from a resource of the computing system **102** via the weight data interface **208** (FIG. 2).

[0179] At block **1510**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the accelerator is to generate at least one of activation sparsity data or weight sparsity data. For example, the configuration circuitry **214** may determine to configure one or more configuration registers of the internal activation sparsity generation circuitry **220** (FIG. 2) and/or the internal weight sparsity generation circuitry **228** (FIG. 2) based on the configuration data **122** of FIG. 1 to generate activation sparsity data and/or weight sparsity data based on the depthwise convolution operation **500**, the first elementwise addition operation **900**, etc.

[0180] If, at block **1510**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the accelerator is not to generate the at least one of the activation sparsity data or the weight sparsity data, control proceeds to block **1514**. If, at block **1510**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the accelerator is to generate the at least one of the activation sparsity data or the weight sparsity data, then, at block **1512**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may generate the at least one of the activation sparsity data or the weight sparsity data based on the acceleration operation. For example, the internal weight sparsity generation circuitry **228** may generate the weight sparsity data to include sparse weight vectors based on a walking is data pattern, such as the third sparsity bit mask vectors **606** of FIG. 6, which may be based on the identification of the acceleration operation as the depthwise convolution operation **500**.

[0181] At block **1514**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may receive the at least one of the activation sparsity data or the weight sparsity data. For example, in response to a determination not to generate the at least one of the activation sparsity data or the weight sparsity data, the external activation sparsity buffer **218** (FIG. 2) may receive activation sparsity data from a resource of the computing system **102** via the activation sparsity data interface **206** (FIG. 2). In some examples, in response to a determination to generate the at least one of the activation sparsity data or the weight sparsity data, the external activation sparsity buffer **218** may receive activation sparsity data from a resource of the computing system **102** via the activation sparsity data interface **206** and/or the external weight sparsity buffer **226** (FIG. 2) may receive weight sparsity data from a resource of the computing system **102** via the weight sparsity data interface **210** (FIG. 2).

[0182] At block **1516**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may generate a sparsity bit mask based on the activation sparsity data and the weight sparsity data. For example, the configuration circuitry **214** may instruct the first multiplexer **230** (FIG. 2) to provide a sparse activation vector to the combined sparsity controller **234** (FIG. 2) based on the acceleration operation. In some examples, the configuration circuitry **214** may instruct the second multiplexer **232** (FIG. 2) to provide a sparse weight vector to the combined sparsity controller **234** based on the acceleration operation. In some such examples, the configuration circuitry **214** may instruct the second multiplexer **232** to select the sparse weight vector from the internal weight

sparsity generation circuitry **228** in response to an identification of the acceleration operation as the depthwise convolution operation **500**. In some such examples, the combined sparsity controller **234** may generate a sparsity bit mask, such as the third sparsity bit mask vectors **606** of FIG. **6**, one of the bit masks **802, 804, 806, 808, 810**, etc., of FIG. **8**, etc., based on the activation sparsity data and the weight sparsity data.

[0183] At block **1518**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may instruct data buffer(s) to provide data to accelerator logic based on the sparsity bit mask. For example, the combined sparsity controller **234** may instruct the activation data buffer **216** to provide activation data based on the sparsity bit mask and the activation sparsity data to the MACs **236**. In some examples, the combined sparsity controller **234** may instruct the weight data buffer **222** to provide weight data based on the sparsity bit mask and the weight sparsity data to the MACs **236**.

[0184] At block **1520**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may execute the acceleration operation with the accelerator logic. For example, the accelerator **200** may execute the depthwise convolution operation **500** of FIG. **5**, the grouped convolution operation **700** of FIG. **7**, the first elementwise addition operation **900** of FIG. **9**, the second elementwise addition operation **1000** of FIG. **10**, the third elementwise addition operation **1100** of FIG. **11**, the first dilated convolution operation **1202** of FIG. **12**, the second dilated convolution operation **1204** of FIG. **12**, the third dilated convolution operation **1206** of FIG. **12**, etc.

[0185] At block **1522**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether to execute another acceleration operation. For example, the configuration circuitry **214** may determine whether to select another acceleration operation to execute with the accelerator **200**. If, at block **1522**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine to execute another acceleration operation, control returns to block **1502** to identify the acceleration operation based on configuration information associated with the accelerator, otherwise the example process **1500** of FIG. **15** concludes.

[0186] FIG. **16** is a flowchart representative of an example process **1600** that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, and/or the accelerator **200** of FIG. **2** to execute an acceleration operation based on an acceleration mode. The example process **1600** begins at block **1602**, at which the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether an acceleration mode is two-dimensional (2-D) convolution. For example, the configuration circuitry **214** (FIG. **2**) may determine that the acceleration mode of the accelerator **200** is to be a 2-D convolution mode based on the configuration data **122** of FIG. **1** received via the configuration interface **202** (FIG. **2**).

[0187] If, at block **1602**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is not 2-D convolution, control proceeds to block **1606**. If, at block **1602**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is 2-D convolution, then, at block **1604**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** obtain sparse weight vector(s) from an acceleration interface. For example, the external weight sparsity buffer **226** (FIG. **2**) may receive weight sparsity data including one or more sparse weight vectors from a resource of the computing system **102** via the weight sparsity data interface **210** (FIG. **2**).

[0188] At block **1606**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the acceleration mode is depthwise convolution. For example, the configuration circuitry **214** may determine that the acceleration mode of the accelerator **200** is to be a depthwise convolution mode based on the configuration data **122**.

[0189] If, at block **1606**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is not depthwise convolution, control proceeds to block **1610**. If, at block **1606**, the first acceleration circuitry **108** and/or the second acceleration circuitry

110 determine that the acceleration mode is depthwise convolution, then, at block **1608**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may identify a sparse weight vector data pattern based on depthwise convolution. For example, the configuration circuitry **214** may instruct the internal weight sparsity generation circuitry **228** (FIG. 2) to generate one or more sparse weight vectors based on a sparse weight vector data pattern. In some such examples, the internal weight sparsity generation circuitry **228** may identify the sparse weight vector data pattern in response to a determination that the accelerator **200** is to execute the depthwise convolution operation **500** of FIG. 5 based on the acceleration mode being depthwise convolution.

[0190] At block **1610**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the acceleration mode is grouped convolution. For example, the configuration circuitry **214** may determine that the acceleration mode of the accelerator **200** is to be a grouped convolution mode based on the configuration data **122**.

[0191] If, at block **1610**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is not grouped convolution, control proceeds to block **1614**. If, at block **1610**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is grouped convolution, then, at block **1612**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may identify a sparse weight vector data pattern based on grouped convolution. For example, the configuration circuitry **214** may instruct the internal weight sparsity generation circuitry **228** to generate one or more sparse weight vectors based on a sparse weight vector data pattern. In some such examples, the internal weight sparsity generation circuitry **228** may identify the sparse weight vector data pattern in response to a determination that the accelerator **200** is to execute the grouped convolution operation **700** of FIG. 7 based on the acceleration mode being grouped convolution. In some such examples, the internal weight sparsity generation circuitry **228** may generate one(s) of the bit masks **802, 804, 806, 808, 810**.

[0192] At block **1614**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the acceleration mode is elementwise addition. For example, the configuration circuitry **214** may determine that the acceleration mode of the accelerator **200** is to be an elementwise addition mode based on the configuration data **122**.

[0193] If, at block **1614**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is not elementwise addition, control proceeds to block **1618**. If, at block **1614**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is elementwise addition, then, at block **1616**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may identify a sparse weight vector data pattern based on elementwise addition. For example, the configuration circuitry **214** may instruct the internal weight sparsity generation circuitry **228** to generate one or more sparse weight vectors based on a sparse weight vector data pattern. In some such examples, the internal weight sparsity generation circuitry **228** may identify the sparse weight vector data pattern in response to a determination that the accelerator **200** is to execute the first elementwise addition operation **900** of FIG. 9, the second elementwise addition operation **1000** of FIG. 10, and/or the third elementwise addition operation **1100** of FIG. 11 based on the acceleration mode being elementwise addition. In some such examples, the internal weight sparsity generation circuitry **228** may generate the columns of weight buffers **906** of FIG. 9.

[0194] At block **1618**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the acceleration mode is dilated convolution. For example, the configuration circuitry **214** may determine that the acceleration mode of the accelerator **200** is to be a dilated convolution mode based on the configuration data **122**.

[0195] If, at block **1618**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is not dilated convolution, control proceeds to block **1624**. If, at block **1618**, the first acceleration circuitry **108** and/or the second acceleration circuitry

110 determine that the acceleration mode is dilated convolution, then, at block **1620**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may identify a sparse weight vector pattern based on dilated convolution. For example, the configuration circuitry **214** may instruct the internal weight sparsity generation circuitry **228** to generate one or more sparse weight vectors based on a sparse weight vector data pattern. In some such examples, the internal weight sparsity generation circuitry **228** may identify the sparse weight vector data pattern in response to a determination that the accelerator **200** is to execute one(s) of the dilated convolution operations **1202**, **1204**, **1206** of FIG. **12** based on the acceleration mode being dilated convolution. In some such examples, the internal weight sparsity generation circuitry **228** may generate the second kernel **1304** of FIG. **13**.

[0196] At block **1622**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may instruct data buffers to provide activation data and weight data based on the sparse weight vector data pattern. For example, the combined sparsity controller **234** (FIG. **2**) may instruct the activation data buffer **216** (FIG. **2**) to provide activation data based on the activation data and a sparsity bit mask, which may be based on sparse activation vectors and sparse weight vectors based on the sparse weight vector data pattern, to the MACs **236** (FIG. **2**). In some examples, the combined sparsity controller **234** may instruct the weight data buffer **222** (FIG. **2**) to provide weight data based on the sparsity bit mask and weight data to the MACs **236**.

[0197] At block **1624**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether the acceleration mode is a custom mode. For example, the configuration circuitry **214** may determine that the acceleration mode of the accelerator **200** is a custom mode in which a user, a computing system, etc., may configure the accelerator **200** for a custom operation, a custom acceleration operation, etc.

[0198] If, at block **1624**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is not the custom mode, control proceeds to block **1628**. If, at block **1624**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine that the acceleration mode is the custom mode, then, at block **1626**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may execute a custom acceleration based on configuration data. For example, the user, the computing system, etc., may provide the configuration circuitry **214** with configuration data including a first instruction to cause the accelerator **200** to internally generate sparse weight vector(s) based on a first data pattern, a second instruction to cause the accelerator **200** to internally generate weight data based on a second data pattern, a third instruction to cause the accelerator **200** to internally generate sparse activation vector(s) based on a third data pattern, etc., and/or a combination thereof. In some such examples, in response to the first instruction, the internal activation sparsity generation circuitry **220** may generate the sparse activation vector(s) based on the first data pattern. In some examples, in response to the second instruction, the internal weight sparsity generation circuitry **228** may generate the sparse weight vector(s) based on the second data pattern. In some examples, in response to the third instruction, the internal weight generation circuitry **224** may generate the weight data based on the third data pattern. In some such examples, the combined sparsity controller **234** may instruct the activation data buffer **216** to provide activation data based on the activation data and a sparsity bit mask, which may be based on sparse activation vector(s) and sparse weight vector(s) based on one(s) of the first data pattern and/or the second data pattern, to the MACs **236**. In some examples, the combined sparsity controller **234** may instruct the weight data buffer **222** to provide weight data based on the sparsity bit mask and the internally generated weight data to the MACs **236**. In some such examples, the MACs **236** may execute the custom operation based on the provided activation and weight data.

[0199] At block **1628**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** may determine whether to select another acceleration operation to execute. For example, the configuration circuitry **214** may determine whether to select another acceleration operation to

execute with the accelerator **200**. If, at block **1628**, the first acceleration circuitry **108** and/or the second acceleration circuitry **110** determine to select another acceleration operation to execute, control returns to block **1602**, otherwise the example process **1600** of FIG. **16** concludes.

[0200] FIG. **17** is a flowchart representative of an example process **1700** that may be performed using example machine readable instructions that may be executed and/or hardware configured to implement the accelerator compiler **104A-C** of FIG. **1** and/or the accelerator compiler **300** of FIG. **3** to generate an executable based on a compiled configuration of a machine-learning model. The example process **1700** of FIG. **17** begins at block **1702**, at which the accelerator compiler **104A-C** may identify a machine-learning model to be executed by an accelerator. For example, the ML model configuration handler **310** (FIG. **3**) may identify one of the ML model(s) **124** of FIG. **1** to be executed by the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry **110** of FIG. **1**, the accelerator **200** of FIG. **2**, etc. In some such examples, the ML model configuration handler **310** may identify the one of the ML model(s) **124** to be a NN to be executed by the accelerator **200**.

[0201] At block **1704**, the accelerator compiler **104A-C** may identify an acceleration mode of the accelerator. For example, the ML model configuration handler **310** may identify a depthwise convolution mode of the accelerator **200** based on the identified NN, or layer(s) thereof to process.

[0202] At block **1706**, the accelerator compiler **104A-C** may determine to provide at least one of activation data or activation sparsity data to the accelerator based on the acceleration mode. For example, the activation configuration handler **320** (FIG. **3**) may determine to provide at least one of activation data or activation sparsity data to the accelerator **200** based on the depthwise convolution mode. In some examples, the activation configuration handler **320** may determine not to provide activation data and/or activation sparsity data to the accelerator **200** based on the depthwise convolution mode to improve bandwidth of the accelerator. In some such examples, the activation configuration handler **320** may determine that the accelerator **200** can internally generate the activation sparsity data to reduce a quantity of data to be read in to the accelerator **200** and thereby improve the efficiency of the accelerator **200**.

[0203] If, at block **1706**, the accelerator compiler **104A-C** determines not to provide at least one of the activation data or the activation sparsity data to the accelerator based on the acceleration mode, control proceeds to block **1710**. If, at block **1706**, the accelerator compiler **104A-C** determines to provide at least one of the activation data or the activation sparsity data to the accelerator based on the acceleration mode, then, at block **1708**, the accelerator compiler **104A-C** may identify the at least one of the activation data or the activation sparsity data to be provided to the accelerator. For example, the activation configuration handler **320** may identify activation data to be transmitted to the activation data interface **204** of FIG. **2** and/or activation sparsity data to be transmitted to the activation sparsity data interface **206** of FIG. **2**.

[0204] At block **1710**, the accelerator compiler **104A-C** may determine whether to provide at least one of weight data or weight sparsity data to the accelerator based on the acceleration mode. For example, the weight configuration handler **330** (FIG. **3**) may determine to provide at least one of weight data or weight sparsity data to the accelerator **200** based on the depthwise convolution mode. In some examples, the weight configuration handler **330** may determine not to provide weight data and/or weight sparsity data to the accelerator **200** based on the depthwise convolution mode to improve bandwidth of the accelerator. In some such examples, the weight configuration handler **330** may determine that the accelerator **200** can internally generate the weight data and/or the weight sparsity data to reduce a quantity of data to be read in to the accelerator **200** and thereby improve the efficiency of the accelerator **200**.

[0205] If, at block **1710**, the accelerator compiler **104A-C** determines not to provide at least one of the weight data or the weight sparsity data to the accelerator based on the acceleration mode, control proceeds to block **1714**. If, at block **1710**, the accelerator compiler **104A-C** determines to provide at least one of the weight data or the weight sparsity data to the accelerator based on the

acceleration mode, then, at block **1712**, the accelerator compiler **104A-C** may identify the at least one of the weight data or the weight sparsity data to be provided to the accelerator. For example, the weight configuration handler **330** may identify weight data to be transmitted to the weight data interface **208** of FIG. 2 and/or weight sparsity data to be transmitted to the weight sparsity data interface **210** of FIG. 2.

[0206] At block **1714**, the accelerator compiler **104A-C** may compile a configuration of the machine-learning model to be executed by the accelerator based on the acceleration mode. For example, the configuration compiler **340** (FIG. 3) may compile a configuration of one or more layers of the NN based on the configuration data **122** of FIG. 1. In some examples, the configuration compiler **340** may compile the configuration of the one or more NN layers in response to a determination not to provide at least one of activation sparsity data, weight sparsity data, or weight data to the accelerator **200**. In some such examples, the configuration compiler **340** may compile the configuration to include a command, an instruction, etc., to the accelerator **200** to generate activation sparsity data with the internal activation sparsity generation circuitry **220**, generate weight sparsity data with the internal weight sparsity generation circuitry **228**, and/or generate weight data with the internal weight generation circuitry **224**.

[0207] At block **1716**, the accelerator compiler **104A-C** may generate an executable based on the compiled configuration. For example, the executable generator **350** (FIG. 3) may generate an executable based on the compiled one(s) of the one or more NN layers. In some such examples, the executable generator **350** may store the executable as one of the ML model(s) **124**. In some such examples, the interface circuitry **114** (FIG. 1) may transmit the executable to one(s) of the external computing systems **130**, which, when executed by the one(s) of the external computing systems **130**, may configure and/or execute a NN associated with the one(s) of the external computing systems **130**.

[0208] At block **1718**, the accelerator compiler **104A-C** may provide the at least one of the activation data, the weight data, the activation sparsity data, or the weight sparsity data to the accelerator. For example, the acceleration interface **360** (FIG. 3) may transmit activation data to the activation data interface **204**, activation sparsity data to the activation sparsity data interface **206**, weight data to the weight data interface **208**, and/or weight sparsity data to the weight sparsity data interface **210** of the accelerator **200**.

[0209] At block **1720**, the accelerator compiler **104A-C** may instruct the accelerator to execute an acceleration operation. For example, the accelerator interface **360** may transmit the compiled configuration(s) to the configuration interface **202** that, when received, may invoke and/or otherwise cause the accelerator **200** to execute the depthwise convolution operation (or a different acceleration operation) based on the compiled configuration(s). In response to instructing the accelerator to execute the acceleration operation at block **1720**, the example process **1700** of FIG. 17 concludes.

[0210] FIG. 18 is a block diagram of an example processor platform **1800** structured to execute the instructions of FIGS. 14-17 to implement the acceleration compiler **104A-C** of FIG. 1, the first acceleration circuitry **108** of FIG. 1, the second acceleration circuitry **110** of FIG. 1, the accelerator **200** of FIG. 2, and/or the accelerator compiler **300** of FIG. 3. The processor platform **1800** can be, for example, a server, a personal computer, a workstation, a self-learning machine (e.g., a neural network), a mobile device (e.g., a cell phone, a smart phone, a tablet such as an iPad™), a personal digital assistant (PDA), an Internet appliance, a DVD player, a CD player, a digital video recorder, a Blu-ray player, a gaming console, a personal video recorder, a set top box, a headset or other wearable device, or any other type of computing device.

[0211] The processor platform **1800** of the illustrated example includes processor circuitry **1812**. The processor circuitry **1812** of the illustrated example is hardware. For example, the processor circuitry **1812** can be implemented by one or more integrated circuits, logic circuits, microprocessors, GPUs, DSPs, or controllers from any desired family or manufacturer. The

hardware processor may be a semiconductor based (e.g., silicon based) device. In this example, the processor circuitry **1812** implements the example ML model configuration handler **310**, the example activation configuration handler **320**, the example weight configuration handler **330**, the example configuration compiler **340**, and the example executable generator **350** of FIG. 3.

[0212] The processor circuitry **1812** of the illustrated example includes a local memory **1813** (e.g., a cache). The processor circuitry **1812** of the illustrated example is in communication with a main memory including a volatile memory **1814** and a non-volatile memory **1816** via a bus **1818**. For example, the bus **1818** may implement the example bus **116** of FIG. 1 and/or the example bus **395** of FIG. 3. The volatile memory **1814** may be implemented by SDRAM, DRAM, RDRAM®, and/or any other type of random access memory device. The non-volatile memory **1816** may be implemented by flash memory and/or any other desired type of memory device. Access to the main memory **1814**, **1816** is controlled by a memory controller **1817**.

[0213] The processor platform **1800** of the illustrated example also includes an interface circuit **1820**. The interface circuit **1820** may be implemented by any type of interface standard, such as an Ethernet interface, a USB, a Bluetooth® interface, an NFC interface, and/or a PCIe interface.

[0214] In the illustrated example, one or more input devices **1822** are connected to the interface circuit **1820**. The input device(s) **1822** permit(s) a user to enter data and/or commands into the processor circuitry **1812**. The input device(s) can be implemented by, for example, an audio sensor, a microphone, a camera (still or video), a keyboard, a button, a mouse, a touchscreen, a track-pad, a trackball, an isopoint device, and/or a voice recognition system.

[0215] One or more output devices **1824** are also connected to the interface circuit **1820** of the illustrated example. The output devices **1824** can be implemented, for example, by display devices (e.g., a light emitting diode (LED), an organic light emitting diode (OLED), a liquid crystal display (LCD), a cathode ray tube (CRT) display, an in-place switching (IPS) display, a touchscreen, etc.), a tactile output device, a printer, and/or speaker. The interface circuit **1820** of the illustrated example, thus, typically includes a graphics driver card, a graphics driver chip and/or a graphics driver processor.

[0216] The interface circuit **1820** of the illustrated example also includes a communication device such as a transmitter, a receiver, a transceiver, a modem, a residential gateway, a wireless access point, and/or a network interface to facilitate exchange of data with external machines (e.g., computing devices of any kind) via a network **1826**. The communication can be via, for example, an Ethernet connection, a DSL connection, a telephone line connection, a coaxial cable system, a satellite system, a line-of-site wireless system, a cellular telephone system, etc. In this example, the interface circuit **1820** implements the example accelerator interface **360** of FIG. 3.

[0217] The processor platform **1800** of the illustrated example also includes one or more mass storage devices **1828** for storing software and/or data. Examples of such mass storage devices **1828** include floppy disk drives, hard drive disks, compact disk drives, Blu-ray disk drives, redundant array of independent disks (RAID) systems, and DVD drives.

[0218] The machine executable instructions **1832** of FIGS. 14-17 may be stored in the mass storage device **1828**, in the volatile memory **1814**, in the non-volatile memory **1816**, and/or on a removable non-transitory computer readable storage medium such as a CD or DVD. In this example, the one or more mass storage devices **1828** implement the example ML model datastore **370**, the example activation datastore **380**, and the example weight datastore **390** of FIG. 3.

[0219] The processor platform **1800** of the illustrated example of FIG. 18 includes an example graphics processing unit (GPU) **1840**, an example vision processing unit (VPU) **1842**, and an example neural network processor **1844**. In this example, the GPU **1840**, the VPU **1842**, and the neural network processor **1844** are in communication with different hardware of the processor platform **1800**, such as the volatile memory **1814**, the non-volatile memory **1816**, etc., via the bus **1818**. In this example, the neural network processor **1844** may be implemented by one or more integrated circuits, logic circuits, microprocessors, GPUs, DSPs, or controllers from any desired

family or manufacturer that can be used to execute an AI model, such as a neural network, which may be implemented by the ML model(s) **124** of FIG. **1**. In some examples, one or more of the ML model configuration handler **310**, the activation configuration handler **320**, the weight configuration handler **330**, the configuration compiler **340**, the executable generator **350**, the accelerator interface **360**, the ML model datastore **370**, the activation datastore **380**, and/or the weight datastore **390** may be implemented in or with at least one of the GPU **1840**, the VPU **1842**, or the neural network processor **1844** instead of or in addition to the processor circuitry **1812**, the one or more mass storage devices **1828**, and/or the interface circuit **1820**.

[0220] In some examples, the graphics processing unit **1840** may implement the first acceleration circuitry **108**, the second acceleration circuitry **110**, and/or the general purpose processing circuitry **112** of FIG. **1**. In some examples, the VPU **1842** may implement the first acceleration circuitry **108**, the second acceleration circuitry **110**, and/or the general purpose processing circuitry **112** of FIG. **1**. In some examples, the neural network processor **1844** may implement the first acceleration circuitry **108**, the second acceleration circuitry **110**, and/or the general purpose processing circuitry **112** of FIG. **1**.

[0221] FIG. **19** is a block diagram of an example implementation of the processor circuitry **1812** of FIG. **18**. In this example, the processor circuitry **1812** of FIG. **18** is implemented by a microprocessor **1900**. For example, the microprocessor **1900** may implement multi-core hardware circuitry such as a CPU, a DSP, a GPU, an XPU, etc. Although it may include any number of example cores **1902** (e.g., **1** core), the microprocessor **1900** of this example is a multi-core semiconductor device including **N** cores. The cores **1902** of the microprocessor **1900** may operate independently or may cooperate to execute machine readable instructions. For example, machine code corresponding to a firmware program, an embedded software program, or a software program may be executed by one of the cores **1902** or may be executed by multiple ones of the cores **1902** at the same or different times. In some examples, the machine code corresponding to the firmware program, the embedded software program, or the software program is split into threads and executed in parallel by two or more of the cores **1902**. The software program may correspond to a portion or all of the machine readable instructions and/or operations represented by the flowcharts of FIGS. **14-17**.

[0222] The cores **1902** may communicate by an example bus **1904**. In some examples, the bus **1904** may implement a communication bus to effectuate communication associated with one(s) of the cores **1902**. For example, the bus **1904** may implement at least one of an Inter-Integrated Circuit (I2C) bus, a Serial Peripheral Interface (SPI) bus, a PCI bus, or a PCIe bus. Additionally or alternatively, the bus **1904** may implement any other type of computing or electrical bus. The cores **1902** may obtain data, instructions, and/or signals from one or more external devices by example interface circuitry **1906**. The cores **1902** may output data, instructions, and/or signals to the one or more external devices by the interface circuitry **1906**. Although the cores **1902** of this example include example local memory **1920** (e.g., Level 1 (L1) cache that may be split into an L1 data cache and an L1 instruction cache), the microprocessor **1900** also includes example shared memory **1910** that may be shared by the cores (e.g., Level 2 (L2_cache)) for high-speed access to data and/or instructions. Data and/or instructions may be transferred (e.g., shared) by writing to and/or reading from the shared memory **1910**. The local memory **1920** of each of the cores **1902** and the shared memory **1910** may be part of a hierarchy of storage devices including multiple levels of cache memory and the main memory (e.g., the main memory **1814**, **1816** of FIG. **18**). Typically, higher levels of memory in the hierarchy exhibit lower access time and have smaller storage capacity than lower levels of memory. Changes in the various levels of the cache hierarchy are managed (e.g., coordinated) by a cache coherency policy.

[0223] Each core **1902** may be referred to as a CPU, DSP, GPU, etc., or any other type of hardware circuitry. Each core **1902** includes control unit circuitry **1914**, arithmetic and logic (AL) circuitry (sometimes referred to as an ALU) **1916**, a plurality of registers **1918**, the L1 cache **1920**, and an

example bus **1922**. Other structures may be present. For example, each core **1902** may include vector unit circuitry, single instruction multiple data (SIMD) unit circuitry, load/store unit (LSU) circuitry, branch/jump unit circuitry, floating-point unit (FPU) circuitry, etc. The control unit circuitry **1914** includes semiconductor-based circuits structured to control (e.g., coordinate) data movement within the corresponding core **1902**. The AL circuitry **1916** includes semiconductor-based circuits structured to perform one or more mathematic and/or logic operations on the data within the corresponding core **1902**. The AL circuitry **1916** of some examples performs integer based operations. In other examples, the AL circuitry **1916** also performs floating point operations. In yet other examples, the AL circuitry **1916** may include first AL circuitry that performs integer based operations and second AL circuitry that performs floating point operations. In some examples, the AL circuitry **1916** may be referred to as an Arithmetic Logic Unit (ALU). The registers **1918** are semiconductor-based structures to store data and/or instructions such as results of one or more of the operations performed by the AL circuitry **1916** of the corresponding core **1902**. For example, the registers **1918** may include vector register(s), SIMD register(s), general purpose register(s), flag register(s), segment register(s), machine specific register(s), instruction pointer register(s), control register(s), debug register(s), memory management register(s), machine check register(s), etc. The registers **1918** may be arranged in a bank as shown in FIG. **19**.

Alternatively, the registers **1918** may be organized in any other arrangement, format, or structure including distributed throughout the core **1902** to shorten access time. The bus **1920** may implement at least one of an I2C bus, a SPI bus, a PCI bus, or a PCIe bus

[0224] Each core **1902** and/or, more generally, the microprocessor **1900** may include additional and/or alternate structures to those shown and described above. For example, one or more clock circuits, one or more power supplies, one or more power gates, one or more cache home agents (CHAs), one or more converged/common mesh stops (CMSs), one or more shifters (e.g., barrel shifter(s)) and/or other circuitry may be present. The microprocessor **1900** is a semiconductor device fabricated to include many transistors interconnected to implement the structures described above in one or more integrated circuits (ICs) contained in one or more packages. The processor circuitry may include and/or cooperate with one or more accelerators. In some examples, accelerators are implemented by logic circuitry to perform certain tasks more quickly and/or efficiently than can be done by a general purpose processor. Examples of accelerators include ASICs and FPGAs such as those discussed herein. A GPU or other programmable device can also be an accelerator. Accelerators may be on-board the processor circuitry, in the same chip package as the processor circuitry and/or in one or more separate packages from the processor circuitry.

[0225] FIG. **20** is a block diagram of another example implementation of the processor circuitry **1812** of FIG. **18**. In this example, the processor circuitry **1812** is implemented by FPGA circuitry **2000**. The FPGA circuitry **2000** can be used, for example, to perform operations that could otherwise be performed by the example microprocessor **1900** of FIG. **19** executing corresponding machine readable instructions. However, once configured, the FPGA circuitry **2000** instantiates the machine readable instructions in hardware and, thus, can often execute the operations faster than they could be performed by a general purpose microprocessor executing the corresponding software.

[0226] More specifically, in contrast to the microprocessor **1900** of FIG. **19** described above (which is a general purpose device that may be programmed to execute some or all of the machine readable instructions represented by the flowchart of FIGS. **14-17** but whose interconnections and logic circuitry are fixed once fabricated), the FPGA circuitry **2000** of the example of FIG. **20** includes interconnections and logic circuitry that may be configured and/or interconnected in different ways after fabrication to instantiate, for example, some or all of the machine readable instructions represented by the flowcharts of FIGS. **14-17**. In particular, the FPGA **2000** may be thought of as an array of logic gates, interconnections, and switches. The switches can be programmed to change how the logic gates are interconnected by the interconnections, effectively

forming one or more dedicated logic circuits (unless and until the FPGA circuitry **2000** is reprogrammed). The configured logic circuits enable the logic gates to cooperate in different ways to perform different operations on data received by input circuitry. Those operations may correspond to some or all of the software represented by the flowcharts of FIGS. **14-17**. As such, the FPGA circuitry **2000** may be structured to effectively instantiate some or all of the machine readable instructions of the flowcharts of FIG. **14-17** as dedicated logic circuits to perform the operations corresponding to those software instructions in a dedicated manner analogous to an ASIC. Therefore, the FPGA circuitry **2000** may perform the operations corresponding to the some or all of the machine readable instructions of FIGS. **14-17** faster than the general purpose microprocessor can execute the same.

[0227] In the example of FIG. **20**, the FPGA circuitry **2000** is structured to be programmed (and/or reprogrammed one or more times) by an end user by a hardware description language (HDL) such as Verilog. The FPGA circuitry **2000** of FIG. **20**, includes example input/output (I/O) circuitry **2002** to obtain and/or output data to/from example configuration circuitry **2004** and/or external hardware (e.g., external hardware circuitry) **2006**. For example, the configuration circuitry **2004** may implement interface circuitry that may obtain machine readable instructions to configure the FPGA circuitry **2000**, or portion(s) thereof. In some such examples, the configuration circuitry **2004** may obtain the machine readable instructions from a user, a machine (e.g., hardware circuitry (e.g., programmed or dedicated circuitry) that may implement an Artificial Intelligence/Machine Learning (AI/ML) model to generate the instructions), etc. In some examples, the external hardware **2006** may implement the microprocessor **1900** of FIG. **19**. In some examples, the external hardware **2006** may implement the first acceleration circuitry **108** of FIG. **1**, the second acceleration circuitry of FIG. **1**, and/or any other component of FIG. **1**. In some examples, the external hardware **2006** may implement the accelerator **200** of FIG. **2**. The FPGA circuitry **2000** also includes an array of example logic gate circuitry **2008**, a plurality of example configurable interconnections **2010**, and example storage circuitry **2012**. The logic gate circuitry **2008** and interconnections **2010** are configurable to instantiate one or more operations that may correspond to at least some of the machine readable instructions of FIGS. **14-17** and/or other desired operations. The logic gate circuitry **2008** shown in FIG. **20** is fabricated in groups or blocks. Each block includes semiconductor-based electrical structures that may be configured into logic circuits. In some examples, the electrical structures include logic gates (e.g., And gates, Or gates, Nor gates, etc.) that provide basic building blocks for logic circuits. Electrically controllable switches (e.g., transistors) are present within each of the logic gate circuitry **2008** to enable configuration of the electrical structures and/or the logic gates to form circuits to perform desired operations. The logic gate circuitry **2008** may include other electrical structures such as look-up tables (LUTs), registers (e.g., flip-flops or latches), multiplexers, etc.

[0228] The interconnections **2010** of the illustrated example are conductive pathways, traces, vias, or the like that may include electrically controllable switches (e.g., transistors) whose state can be changed by programming (e.g., using an HDL instruction language) to activate or deactivate one or more connections between one or more of the logic gate circuitry **2008** to program desired logic circuits.

[0229] The storage circuitry **2012** of the illustrated example is structured to store result(s) of the one or more of the operations performed by corresponding logic gates. The storage circuitry **2012** may be implemented by registers or the like. In the illustrated example, the storage circuitry **2012** is distributed amongst the logic gate circuitry **2008** to facilitate access and increase execution speed.

[0230] The example FPGA circuitry **2000** of FIG. **20** also includes example Dedicated Operations Circuitry **2014**. In this example, the Dedicated Operations Circuitry **2014** includes special purpose circuitry **2016** that may be invoked to implement commonly used functions to avoid the need to program those functions in the field. Examples of such special purpose circuitry **2016** include memory (e.g., DRAM) controller circuitry, PCIe controller circuitry, clock circuitry, transceiver

circuitry, memory, and multiplier-accumulator circuitry. Other types of special purpose circuitry may be present. In some examples, the FPGA circuitry **2000** may also include example general purpose programmable circuitry **2018** such as an example CPU **2020** and/or an example DSP **2022**. Other general purpose programmable circuitry **2018** may additionally or alternatively be present such as a GPU, an XPU, etc., that can be programmed to perform other operations.

[0231] Although FIGS. **19** and **20** illustrate two example implementations of the processor circuitry **1812** of FIG. **18**, many other approaches are contemplated. For example, as mentioned above, modern FPGA circuitry may include an on-board CPU, such as one or more of the example CPU **2020** of FIG. **20**. Therefore, the processor circuitry **1812** of FIG. **18** may additionally be implemented by combining the example microprocessor **1900** of FIG. **19** and the example FPGA circuitry **2000** of FIG. **20**. In some such hybrid examples, a first portion of the machine readable instructions represented by the flowcharts of FIG. **14-17** may be executed by one or more of the cores **1902** of FIG. **19** and a second portion of the machine readable instructions represented by the flowcharts of FIGS. **14-17** may be executed by the FPGA circuitry **2000** of FIG. **20**.

[0232] In some examples, the processor circuitry **1812** of FIG. **18** may be in one or more packages. For example, the processor circuitry **1900** of FIG. **19** and/or the FPGA circuitry **2000** of FIG. **20** may be in one or more packages. In some examples, an XPU may be implemented by the processor circuitry **1812** of FIG. **18**, which may be in one or more packages. For example, the XPU may include a CPU in one package, a DSP in another package, a GPU in yet another package, and an FPGA in still yet another package.

[0233] FIG. **21** is a block diagram of an example software distribution platform **2105** to distribute software (e.g., software corresponding to the example machine readable instructions **1832** of FIG. **18**) to third parties, client devices, etc., such as consumers (e.g., for license, sale and/or use), retailers (e.g., for sale, re-sale, license, and/or sub-license), and/or original equipment manufacturers (OEMs) (e.g., for inclusion in products to be distributed to, for example, retailers and/or to direct buy customers). The example software distribution platform **2105** may be implemented by any computer server, data facility, cloud service, etc., capable of storing and transmitting software to other computing devices. The third parties may be customers of the entity owning and/or operating the software distribution platform. For example, the entity that owns and/or operates the software distribution platform may be a developer, a seller, and/or a licensor of software such as the example machine readable instructions **1832** of FIG. **18**. The third parties may be consumers, users, retailers, OEMs, etc., who purchase and/or license the software for use and/or re-sale and/or sub-licensing. In the illustrated example, the software distribution platform **2105** includes one or more servers and one or more storage devices. The storage devices store the machine readable instructions **1832**, which may correspond to the example machine readable instructions **1400**, **1500**, **1600**, **1700** of FIGS. **14-17**, as described above. The one or more servers of the example software distribution platform **2105** are in communication with a network **2110**, which may correspond to any one or more of the Internet and/or any of the example networks **128**, **1826** described above. In some examples, the one or more servers are responsive to requests to transmit the software to a requesting party as part of a commercial transaction. Payment for the delivery, sale and/or license of the software may be handled by the one or more servers of the software distribution platform and/or via a third party payment entity. The servers enable purchasers and/or licensors to download the machine readable instructions **1832** from the software distribution platform **2105**. For example, the software, which may correspond to the example machine readable instructions **1400**, **1500**, **1600**, **1700** of FIGS. **14-17**, may be downloaded to the example processor platform **1800**, which is to execute the machine readable instructions **1832** to implement the example acceleration compiler **104A-C** of FIG. **1**, the first example acceleration circuitry **108** of FIG. **1**, the second example acceleration circuitry **110** of FIG. **1**, the example accelerator **200** of FIG. **1**, and/or the example accelerator compiler **300** of FIG. **3**. In some example, one or more servers of the software distribution platform **2105** periodically offer,

transmit, and/or force updates to the software (e.g., the example machine readable instructions **1832** of FIG. **18**) to ensure improvements, patches, updates, etc. are distributed and applied to the software at the end user devices.

[0234] From the foregoing, it will be appreciated that example systems, methods, apparatus, and articles of manufacture have been disclosed to perform machine-learning model operations on sparse accelerators. The disclosed systems, methods, apparatus, and articles of manufacture utilize bit mask vectors as compute vectors in addition to being sparsity vectors to isolate single activations and/or single channels or groups of channels within a tensor. The disclosed systems, methods, apparatus, and articles of manufacture manipulate bit mask vectors internal to a sparse accelerator to save the bandwidth of reading the bit mask vectors from external memory. The disclosed systems, methods, apparatus, and articles of manufacture alternatively may perform the manipulation offline, load to memory, and read the manipulated bit mask vectors into the sparse accelerator when processing the layer. The disclosed systems, methods, apparatus, and articles of manufacture improve the efficiency of using a computing device by reducing a quantity of data to be read into an accelerator to execute a machine-learning layer and execute non-conventional convolution operations as conventional convolution operations to improve system efficiency and performance. The disclosed methods, apparatus and articles of manufacture are accordingly directed to one or more improvement(s) in the functioning of a computer.

[0235] Example methods, apparatus, systems, and articles of manufacture to perform machine-learning model operations on sparse accelerators are disclosed herein. Further examples and combinations thereof include the following:

[0236] Example 1 includes an apparatus to execute an acceleration operation, the apparatus comprising first circuitry, second circuitry to generate sparsity data based on the acceleration operation, and third circuitry to instruct one or more data buffers to provide at least one of activation data or weight data based on the sparsity data to the first circuitry, the first circuitry to execute the acceleration operation based on the at least one of the activation data or the weight data.

[0237] Example 2 includes the apparatus of example 1, wherein the sparsity data includes a sparse weight vector, and further including fourth circuitry to identify the acceleration operation based on configuration information, the sparse weight vector generated based on the configuration information, and the third circuitry to generate a sparsity bit mask based on a sparse activation vector and the sparse weight vector in response to obtaining the sparse activation vector.

[0238] Example 3 includes the apparatus of example 1, wherein the sparsity data includes at least one of activation sparsity data or weight sparsity data, the one or more data buffers include a weight data buffer and an activation data buffer, and the third circuitry is to generate a combined sparsity bit mask based on the activation sparsity data and the weight sparsity data, instruct the activation data buffer to provide activation data to the first circuitry based on the combined sparsity bit mask, and in response to determining to generate weight data, instruct the weight data buffer to provide the generated weight data to the first circuitry based on the combined sparsity bit mask.

[0239] Example 4 includes the apparatus of example 1, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a depthwise convolution operation, and wherein the second circuitry is to identify a data pattern based on the depthwise convolution operation, and generate the sparse weight vector based on the data pattern, and the third circuitry is to generate a sparsity bit mask based on the sparse weight vector.

[0240] Example 5 includes the apparatus of example 1, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a grouped convolution operation, and wherein the second circuitry is to identify a data pattern based on the grouped convolution operation, the data pattern including a set of logic ones, and generate the sparse weight vector based on the data pattern, and the third circuitry is to generate a sparsity bit mask based on the sparse weight vector.

[0241] Example 6 includes the apparatus of example 1, wherein the acceleration operation is an

elementwise addition operation, and wherein the second circuitry is to identify a data pattern based on the elementwise addition operation, the data pattern based on a walking ones data pattern, and generate a plurality of sparse weight vectors based on the data pattern, and the first circuitry is to store first data values associated with a first tensor based on the plurality of the sparse weight vectors and a first set of the activation data, store second data values associated with a second tensor based on the plurality of the sparse weight vectors and a second set of the activation data, and generate output data values based on a sum of the first data values and the second data values.

[0242] Example 7 includes the apparatus of example 1, wherein the sparsity data includes one or more sparse weight vectors, the acceleration operation is a dilated convolution, and wherein the second circuitry is to, in response to an identification of a data pattern based on the dilated convolution, generate the one or more sparse weight vectors based on the data pattern, the data pattern based on a sparse kernel, and the third circuitry is to generate one or more sparsity bit masks based on the one or more sparse weight vectors.

[0243] Example 8 includes a non-transitory computer readable medium comprising instructions that, when executed, cause at least one accelerator to at least generate, with the at least one accelerator, sparsity data based on an acceleration operation to be executed by the at least one accelerator, instruct one or more data buffers to provide at least one of activation data or weight data based on the sparsity data to accelerator logic of the at least one accelerator, and execute the acceleration operation based on the at least one of the activation data or the weight data.

[0244] Example 9 includes the non-transitory computer readable medium of example 8, wherein the sparsity data includes a sparse weight vector, and the instructions, when executed, cause the at least one accelerator to identify the acceleration operation based on configuration information, the sparse weight vector generated based on the configuration information, and in response to obtaining a sparse activation vector, generate a sparsity bit mask based on the sparse activation vector and the sparse weight vector.

[0245] Example 10 includes the non-transitory computer readable medium of example 8, wherein the sparsity data includes at least one of activation sparsity data or weight sparsity data, the one or more data buffers include a weight data buffer and an activation data buffer, and the instructions, when executed, cause the at least one accelerator to generate a combined sparsity bit mask based on the activation sparsity data and the weight sparsity data, instruct the activation data buffer to provide activation data to the accelerator logic based on the combined sparsity bit mask, and in response to determining to generate weight data, instruct the weight data buffer to provide the generated weight data to the accelerator logic based on the combined sparsity bit mask.

[0246] Example 11 includes the non-transitory computer readable medium of example 8, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a depthwise convolution operation, and the instructions, when executed, cause the at least one accelerator to identify a data pattern based on the depthwise convolution operation, generate the sparse weight vector based on the data pattern, and generate a sparsity bit mask based on the sparse weight vector.

[0247] Example 12 includes the non-transitory computer readable medium of example 8, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a grouped convolution operation, and the instructions, when executed, cause the at least one accelerator to identify a data pattern based on the grouped convolution operation, the data pattern including a set of logic ones, generate the sparse weight vector based on the data pattern, and generate a sparsity bit mask based on the sparse weight vector.

[0248] Example 13 includes the non-transitory computer readable medium of example 8, wherein the acceleration operation is an elementwise addition operation, and the instructions, when executed, cause the at least one accelerator to identify a data pattern based on the elementwise addition operation, the data pattern based on a walking ones data pattern, generate a plurality of sparse weight vectors based on the data pattern, store first data values associated with a first tensor based on the plurality of the sparse weight vectors and a first set of the activation data, store second

data values associated with a second tensor based on the plurality of the sparse weight vectors and a second set of the activation data, and generate output data values based on a sum of the first data values and the second data values.

[0249] Example 14 includes the non-transitory computer readable medium of example 8, wherein the sparsity data includes one or more sparse weight vectors, the acceleration operation is a dilated convolution, and the instructions, when executed, cause the at least one accelerator to identify a data pattern based on the dilated convolution, the data pattern based on a sparse kernel, generate the one or more sparse weight vectors based on the data pattern, and generate one or more sparsity bit masks based on the one or more sparse weight vectors.

[0250] Example 15 includes a method comprising generating, with an accelerator, sparsity data based on an acceleration operation to be executed by the accelerator, instructing one or more data buffers to provide at least one of activation data or weight data based on the sparsity data to accelerator logic of the accelerator, and executing the acceleration operation based on the at least one of the activation data or the weight data.

[0251] Example 16 includes the method of example 15, wherein the sparsity data includes a sparse weight vector, and further including identifying the acceleration operation based on configuration information, the sparse weight vector generated based on the configuration information, and in response to obtaining a sparse activation vector, generating a sparsity bit mask based on the sparse activation vector and the sparse weight vector.

[0252] Example 17 includes the method of example 15, wherein the sparsity data includes at least one of activation sparsity data or weight sparsity data, the one or more data buffers include a weight data buffer and an activation data buffer, and further including generating a combined sparsity bit mask based on the activation sparsity data and the weight sparsity data, instructing the activation data buffer to provide activation data to the accelerator logic based on the combined sparsity bit mask, and in response to determining to generate weight data, instructing the weight data buffer to provide the generated weight data to the accelerator logic based on the combined sparsity bit mask.

[0253] Example 18 includes the method of example 15, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a depthwise convolution operation, and further including identifying a data pattern based on the depthwise convolution operation, generating the sparse weight vector based on the data pattern, and generating a sparsity bit mask based on the sparse weight vector.

[0254] Example 19 includes the method of example 15, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a grouped convolution operation, and further including identifying a data pattern based on the grouped convolution operation, the data pattern including a set of logic ones, generating the sparse weight vector based on the data pattern, and generating a sparsity bit mask based on the sparse weight vector.

[0255] Example 20 includes the method of example 15, wherein the acceleration operation is an elementwise addition operation, and further including identifying a data pattern based on the elementwise addition operation, the data pattern based on a walking ones data pattern, generating a plurality of sparse weight vectors based on the data pattern, storing first data values associated with a first tensor based on the plurality of the sparse weight vectors and a first set of the activation data, storing second data values associated with a second tensor based on the plurality of the sparse weight vectors and a second set of the activation data, and generating output data values based on a sum of the first data values and the second data values.

[0256] Example 21 includes the method of example 15, wherein the sparsity data includes one or more sparse weight vectors, the acceleration operation is a dilated convolution, and further including identifying a data pattern based on the dilated convolution, the data pattern based on a sparse kernel, generating the one or more sparse weight vectors based on the data pattern, and generating one or more sparsity bit masks based on the one or more sparse weight vectors.

[0257] Example 22 includes an apparatus to execute an acceleration operation, the apparatus comprising means for generating sparsity data based on the acceleration operation, means for instructing one or more data buffers to provide at least one of activation data or weight data based on the sparsity data, and means for executing the acceleration operation based on the at least one of the activation data or the weight data.

[0258] Example 23 includes the apparatus of example 22, wherein the sparsity data includes a sparse weight vector, and further including means for identifying the acceleration operation based on configuration information, the sparse weight vector generated based on the configuration information, and the means for generating to generate a sparsity bit mask based on a sparse activation vector and the sparse weight vector in response to obtaining the sparse activation vector.

[0259] Example 24 includes the apparatus of example 22, wherein the sparsity data includes at least one of activation sparsity data or weight sparsity data, the one or more data buffers include a weight data buffer and an activation data buffer, and the means for instructing is to generate a combined sparsity bit mask based on the activation sparsity data and the weight sparsity data, instruct the activation data buffer to provide activation data to the means for executing based on the combined sparsity bit mask, and in response to determining to generate weight data, instruct the weight data buffer to provide the generated weight data to the means for executing based on the combined sparsity bit mask.

[0260] Example 25 includes the apparatus of example 22, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a depthwise convolution operation, and wherein the means for generating is to identify a data pattern based on the depthwise convolution operation, and generate the sparse weight vector based on the data pattern, and the means for instructing is to generate a sparsity bit mask based on the sparse weight vector.

[0261] Example 26 includes the apparatus of example 22, wherein the sparsity data includes a sparse weight vector, the acceleration operation is a grouped convolution operation, and wherein the means for generating is to identify a data pattern based on the grouped convolution operation, the data pattern including a set of logic ones, and generate the sparse weight vector based on the data pattern, and the means for instructing is to generate a sparsity bit mask based on the sparse weight vector.

[0262] Example 27 includes the apparatus of example 22, wherein the acceleration operation is an elementwise addition operation, and wherein the means for generating is to identify a data pattern based on the elementwise addition operation, the data pattern based on a walking ones data pattern, and generate a plurality of sparse weight vectors based on the data pattern, and the means for executing is to store first data values associated with a first tensor based on the plurality of the sparse weight vectors and a first set of the activation data, store second data values associated with a second tensor based on the plurality of the sparse weight vectors and a second set of the activation data, and generate output data values based on a sum of the first data values and the second data values.

[0263] Example 28 includes the apparatus of example 22, wherein the sparsity data includes one or more sparse weight vectors, the acceleration operation is a dilated convolution, and wherein the means for generating is to identify a data pattern based on the dilated convolution, the data pattern based on a sparse kernel, and generate the one or more sparse weight vectors based on the data pattern, and the means for instructing is to generate one or more sparsity bit masks based on the one or more sparse weight vectors.

[0264] Example 29 includes an apparatus comprising internal weight sparsity generation circuitry with a first output, first multiplexer circuitry with a first multiplexer input and a first multiplexer output, the first multiplexer input coupled to the first output, controller circuitry with a first controller input, a first controller output, and a second controller output, the first controller input coupled to the first multiplexer output, and acceleration circuitry having acceleration inputs.

[0265] Example 30 includes the apparatus of example 29, further including configuration circuitry

with a first configuration output, a second configuration output, a third configuration output, and a fourth configuration output, the first configuration output coupled to a first input of the internal weight sparsity generation circuitry, the second configuration output coupled to a first multiplexer control input of the first multiplexer circuitry, internal activation sparsity generation circuitry with a second input and a second output, the second input coupled to the third configuration output, and second multiplexer circuitry with a second multiplexer control input, a second multiplexer input, and a second multiplexer output, the second multiplexer control input coupled to the fourth configuration output, the second multiplexer input coupled to the second output.

[0266] Example 31 includes the apparatus of example 30, further including a configuration interface coupled to a configuration input of the configuration circuitry.

[0267] Example 32 includes the apparatus of example 30, further including an external weight sparsity buffer having a first buffer input and a first buffer output, the first buffer output coupled to a third multiplexer input of the first multiplexer circuitry, an external activation sparsity buffer having a second buffer input and a second buffer output, the second buffer output coupled to a fourth multiplexer input of the second multiplexer circuitry, a weight sparsity data interface coupled to the first buffer input, and an activation sparsity data interface coupled to the second buffer input.

[0268] Example 33 includes the apparatus of example 30, wherein the controller circuitry has a second controller input coupled to the second multiplexer output, and further including an activation data buffer having a first data buffer input and first data buffer outputs, the first data buffer input coupled the first controller output, the first data buffer outputs coupled to respective ones of the accelerator inputs, and a weight data buffer having a second data buffer input and second data buffer outputs, the second data buffer input coupled to the second controller output, the second data buffer outputs coupled to respective ones of the accelerator inputs.

[0269] Example 34 includes the apparatus of example 33, wherein the weight data buffer includes internal weight generation circuitry.

[0270] Example 35 includes the apparatus of example 33, further including an accumulator storage with storage inputs and storage outputs, the storage inputs coupled to respective ones of accelerator outputs of the accelerator circuitry, and an output activation interface coupled to the storage outputs.

[0271] Although certain example systems, methods, apparatus, and articles of manufacture have been disclosed herein, the scope of coverage of this patent is not limited thereto. On the contrary, this patent covers all systems, methods, apparatus, and articles of manufacture fairly falling within the scope of the claims of this patent.

[0272] The following claims are hereby incorporated into this Detailed Description by this reference, with each claim standing on its own as a separate embodiment of the present disclosure.

Claims

1. An apparatus, comprising: one or more multiply accumulators; a mask generator to generate a mask vector; a mask buffer to store a pre-generated mask vector; a buffer to store input data of an operation in a neural network; and a controller to: select one or more data elements from the input data by using the mask vector or the pre-generated mask vector, and transmit the one or more data elements from the buffer to the one or more multiply accumulators, wherein the one or more multiply accumulators are to perform the operation in the neural network by using the one or more data elements.
2. The apparatus of claim 1, further comprising: a multiplexer to: select a mask vector for the operation from the mask vector generated by the mask generator and the pre-generated mask vector stored in the mask buffer, and provide the mask vector for the operation to the controller.
3. The apparatus of claim 2, wherein the multiplexer is further to receive configuration data,

wherein the multiplexer is to select the mask vector for the operation based on the configuration data.

4. The apparatus of claim 3, wherein the configuration data indicates a type of the operation in the neural network.

5. The apparatus of claim 4, wherein the type of the operation in the neural network is a depthwise convolution operation, and the multiplexer is to select the mask vector generated by the mask generator as the mask vector for the operation.

6. The apparatus of claim 1, further comprising: an input data generator to generate additional input data for the operation in the neural network, wherein the one or more multiply accumulators are to perform the operation in the neural network by using the one or more data elements and the additional input data.

7. The apparatus of claim 6, further comprising: another buffer to store the additional input data, wherein the buffer is to store activations for a convolution operation, and the additional buffer is to store weights for the convolution operation.

8. The apparatus of claim 1, further comprising: an additional mask generator to generate an additional mask vector; and an additional mask buffer to store an additional pre-generated mask vector, wherein the controller is to select the one or more data elements from the input data by further using the additional mask vector or the additional pre-generated mask vector.

9. The apparatus of claim 1, wherein the mask vector is a bit mask vector comprising a plurality of bits.

10. The apparatus of claim 9, wherein a one-bit in the mask vector indicates a selection of a data element in the input data for performing the operation.

11. A method, comprising: receiving configuration data indicating a type of an operation in a neural network; selecting, using the configuration data, a mask vector for the operation from a signal from a mask generator and a signal from a mask buffer, the mask generator to generate a mask vector, the mask buffer to store a pre-generated mask vector; selecting, by using the mask vector for the operation, one or more data elements from input data of the operation; transmit the one or more data elements to one or more multiply accumulators; and performing, by the one or more multiply accumulators using the one or more data elements, the operation in the neural network.

12. The method of claim 11, further comprising: generating additional input data for the operation in the neural network; and transmit the additional input data to the one or more multiply accumulators, wherein the one or more multiply accumulators are to perform the operation in the neural network by using the one or more data elements and the additional input data.

13. The method of claim 12, further comprising: storing the input data in a first buffer; and storing the additional input data in a second buffer, wherein the first buffer is to store activations for a convolution operation, and the second buffer is to store weights for the convolution operation.

14. The method of claim 11, further comprising: selecting, using the configuration data, an additional mask vector for the operation from a signal from an additional mask generator and a signal from an additional mask buffer, the additional mask generator to generate an additional mask vector, the additional mask buffer to store an additional pre-generated mask vector.

15. The method of claim 14, wherein the one or more data elements are selected from the input data of the operation by further using the additional mask vector.

16. One or more non-transitory computer-readable media storing instructions executable to perform operations, the operations comprising: receiving configuration data indicating a type of an operation in a neural network; selecting, using the configuration data, a mask vector for the operation from a signal from a mask generator and a signal from a mask buffer, the mask generator to generate a mask vector, the mask buffer to store a pre-generated mask vector; selecting, by using the mask vector for the operation, one or more data elements from input data of the operation; transmit the one or more data elements to one or more multiply accumulators; and performing, by the one or more multiply accumulators using the one or more data elements, the operation in the

neural network.

17. The one or more non-transitory computer-readable media of claim 16, wherein the operations further comprise: generating additional input data for the operation in the neural network; and transmit the additional input data to the one or more multiply accumulators, wherein the one or more multiply accumulators are to perform the operation in the neural network by using the one or more data elements and the additional input data.

18. The one or more non-transitory computer-readable media of claim 17, wherein the operations further comprise: storing the input data in a first buffer; and storing the additional input data in a second buffer, wherein the first buffer is to store activations for a convolution operation, and the second buffer is to store weights for the convolution operation.

19. The one or more non-transitory computer-readable media of claim 16, wherein the operations further comprise: selecting, using the configuration data, an additional mask vector for the operation from a signal from an additional mask generator and a signal from an additional mask buffer, the additional mask generator to generate an additional mask vector, the additional mask buffer to store an additional pre-generated mask vector.

20. The one or more non-transitory computer-readable media of claim 19, wherein the one or more data elements are selected from the input data of the operation by further using the additional mask vector.
