

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265471

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

PAKDAMAN NAEINI; Mahdi et al.

REINFORCEMENT LEARNING FOR REFINEMENT MODELS

Abstract

The disclosed concepts relate to employing a refinement model to refine actions generated by a first machine learning model. In some cases, the first machine learning model can have fixed parameters that are not readily available to be tuned for a new task domain. To overcome this issue, a refinement model can be employed to refine actions output by the first machine learning model. Then, a reward model can be employed to select either first actions output by the first machine learning model or refined actions output by the refinement model.

Inventors: PAKDAMAN NAEINI; Mahdi (Kirkland, WA), DU; Liang (Redmond, WA), WANG; Wenhan (Bellevue, WA), BARADARAN HASHEMI; Homa (Kirkland, WA)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 1000007712392

Assignee: Microsoft Technology Licensing, LLC (Redmond, WA)

Appl. No.: 18/582354

Filed: February 20, 2024

Publication Classification

Int. Cl.: G06N3/092 (20230101); G06N3/0455 (20230101)

U.S. Cl.:

CPC G06N3/092 (20230101); G06N3/0455 (20230101);

Background/Summary

BACKGROUND

[0001] In recent years, generative language models have demonstrated tremendous capability at generating natural language text. For instance, generative language models can summarize existing documents, help users draft new documents, and conduct natural language conversations with users at a very high level. Given adequate training data, a generative language model can learn to be adept at many different language-generating tasks. However, generative language models can be extremely large, e.g., having billions of parameters. As a consequence, the initial training of a generative language model tends to require massive amounts of training data and associated computational resources.

[0002] Once a generative language model is trained, it can be tuned to new task domains using relatively few training examples. However, in some cases, the parameters of a given generative language model are not available to end users. Thus, it is not always possible to tune a generative language model for a particular task. This can limit the utility of a generative language model for task domains for which the generative language model has not been trained.

SUMMARY

[0003] This Summary is provided to introduce a selection of concepts in a simplified form that are further described below in the Detailed Description. This Summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used to limit the scope of the claimed subject matter.

[0004] The description generally relates to techniques for employing a refinement model in a reinforcement learning context. One example includes a method or technique that can be performed on a computing device. The method or technique can include obtaining a state of an environment. The method or technique can also include inputting the state to a first machine learning model. The method or technique can also include receiving a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state. The method or technique can also include inputting the state and the first action to a refinement model. The method or technique can also include receiving a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action. The method or technique can also include providing the state, the first action, and the refined action to a reward model. The method or technique can also include receiving, from the reward model, respective reward values for the first action and the refined action based at least on the reward values, selecting the first action or the refined action as a selected action, and outputting the selected action.

[0005] Another example entails a system that includes a hardware processing unit and a storage resource storing computer-readable instructions. When executed by the hardware processing unit, the computer-readable instructions can cause the system to obtain a state of an environment. The computer-readable instructions can also cause the system to input the state to a first machine learning model. The computer-readable instructions can also cause the system to receive a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state. The computer-readable instructions can also cause the system to input the state and the first action to a refinement model. The computer-readable instructions can also cause the system to receive a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action. The computer-readable instructions can also cause the system to provide the state, the first action, and the refined action to a reward model. The computer-readable instructions can also cause the system to receive, from the reward model, respective reward values for the first action and the refined action. The computer-readable instructions can also cause the system to based at least on the reward values, select the first action or the refined action as a selected action. The computer-readable instructions can also cause the system to output the selected action.

[0006] Another example includes a computer-readable storage medium storing computer-readable instructions which, when executed by a processing unit, cause the processing unit to perform acts. The acts can include obtaining a state of an environment. The acts can also include inputting the state to a first machine learning model. The acts can also include receiving a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state. The acts can also include inputting the state and the first action to a refinement model. The acts can also include receiving a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action. The acts can also include providing the state, the first action, and the refined action to a reward model. The acts can also include receiving, from the reward model, respective reward values for the first action and the refined action. The acts can also include based at least on the reward values, selecting the first action or the refined action as a selected action. The acts can also include outputting the selected action.

[0007] The above-listed examples are intended to provide a quick reference to aid the reader and are not intended to define the scope of the concepts described herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0008] The Detailed Description is described with reference to the accompanying figures. In the figures, the left-most digit(s) of a reference number identifies the figure in which the reference number first appears. The use of similar reference numbers in different instances in the description and the figures may indicate similar or identical items.

[0009] FIG. 1 illustrates an example of a generative language model, consistent with some implementations of the present concepts.

[0010] FIG. 2 illustrates an example of a reinforcement learning scenario, consistent with some implementations of the present concepts.

[0011] FIG. 3 shows an example of a workflow for refining actions, consistent with some implementations of the present concepts.

[0012] FIG. 4A shows an example of a reward model, consistent with some implementations of the present concepts.

[0013] FIG. 4B shows an example of a refinement model, consistent with some implementations of the present concepts.

[0014] FIG. 5 shows an example prompt for a first machine learning model, consistent with some implementations of the present concepts.

[0015] FIG. 6 shows an example user experience, consistent with some implementations of the present concepts.

[0016] FIG. 7 illustrates an example system, consistent with some implementations of the present concepts.

[0017] FIG. 8 is a flow chart of an example method or technique, consistent with some implementations of the present concepts.

DETAILED DESCRIPTION

Overview

[0018] Recent advances in language modeling, such as the development of transformer-based generative language models (e.g., one or more versions of models such as GPT, BLOOM, PaLM, and/or LLAMA), have enabled language models to perform complex tasks for users. For instance, generative language models perform well at tasks such as engaging in dialogs with users, summarizing documents for users, etc. Some generative language models have even achieved milestones such as passing the bar exam.

[0019] Generative language models can learn to perform these complex tasks by being exposed to training data relating to a wide range of concepts. As a consequence, massive amounts of training data are generally involved in training a generative language model. In order to effectively represent the knowledge obtained from the training data, generative language models tend to be extremely large, having billions or trillions of parameters.

[0020] In part because of the expense involved in training a large generative language model, some entities may provide generative language models for end users to employ for inference, while preventing the end users from modifying the parameters of the generative language model. Thus, it is not plausible for the end users to tune these generative language model for their own task domains. As a consequence, the utility of a generative language model can be limited to the training that has already been performed on the model.

[0021] Using the disclosed techniques, a refinement model can be employed to refine the actions generated by first machine learning model. For instance, the first machine learning model can be a generative language model with fixed (e.g., inaccessible) parameters, and the refinement model can be another generative language model with tunable parameters. A reward model can be employed to select between actions generated by the first machine learning model and refined actions generated by the refinement model. Reinforcement learning approaches with human/machine-learning feedback can be employed to train and/or tune the refinement model to refine the actions produced by the first machine learning model. In this manner, the use of the refinement model with the reward model can overcome the limitations of generative language models with fixed parameters.

Machine Learning Paradigms

[0022] There are various types of machine learning models that can be trained to perform a given task. Support vector machines, decision trees, neural networks, and contextual bandits are just a few examples of machine learning frameworks that have been used in a wide variety of applications, such as image processing, natural language processing, etc. Generally, machine learning can involve exposing a model to a training signal and then adapting parameters of the model based on the training signal.

[0023] One way to train a machine learning model involves supervised learning, where a model is trained using labeled training data as a training signal. For instance, the training data can include training examples that have been labeled by a human being or other trusted annotator, and the model can be trained by attempting to predict the labels and adjusting model parameters when the predictions are incorrect. Another approach is unsupervised learning, where a model is trained to learn patterns from unlabeled training data, such as learning by predicting masked tokens from a corpus of documents. In semi-supervised learning, a model is trained using both labeled and unlabeled training data, e.g., by pretraining the model using unsupervised learning and then tuning the pretrained model using labeled training data for a particular task. In reinforcement learning, the model is trained using a reward function, where the model receives a reward for reaching certain specified states.

Neural Networks

[0024] In a neural network, nodes are connected to one another via one or more edges. A neural network can include an input layer, an output layer, and one or more intermediate layers. Individual nodes can process their respective inputs according to a predefined function, and provide an output to a subsequent layer, or, in some cases, a previous layer. The inputs to a given node can be multiplied by a corresponding weight value for an edge between the input and the node. In addition, nodes can have individual bias values that are also used to produce outputs.

[0025] Various training procedures can be applied to learn the edge weights and/or bias values. Neural networks can be trained using supervised learning, semi-supervised learning, unsupervised learning, and/or reinforcement learning. Neural networks can be employed for a very wide range of machine learning applications, such as regression, classification, image generation, natural

language generation, etc.

Generative Models

[0026] A generative model is a machine learning model employed to generate new content.

Generative models can be trained to predict items in sequences of training data. When employed in inference mode, the output of a generative model can include new sequences of items that the model generates. A “generative language model” is a model trained from one or more sources of natural language training data to predict a sequence of output tokens given one or more input tokens. A generative language model can generate new sequences of text given some input prompt, e.g., a query potentially with some additional context. For instance, a generative language model can be implemented as a neural network, e.g., a decoder-based generative language model such as GPT, BLOOM, PaLM, and/or LLAMA or variants thereof, a long short-term memory model, etc. A “large” generative language model is a generative language model with one billion or more parameters.

[0027] In some cases, a generative model can be multi-modal. For instance, in addition to textual inputs and/or outputs, the model may be capable of using images, audio, application states, code, or other modalities as inputs and/or generating images, audio, application states, or code or other modalities as outputs. Note that the term “generative language model” encompasses multi-modal generative models where at least one mode of output includes natural language tokens.

[0028] The term “prompt,” as used herein, refers to input text provided to a generative language model that the generative language model uses to generate output text. A prompt can include a query, e.g., a request for information from the generative language model. A prompt can also include context, or additional information that the generative language model uses to respond to the query. In some cases, a prompt can include one or more examples for the generative language model as context (e.g., “few-shot prompting”), and can condition the generative language model to generate more accurate responses than the generative model would produce without the examples. The term “in-context learning,” as used herein, refers to learning, by a generative model, from examples input to the model at inference time, where the examples enable the generative model to learn without performing explicit training, e.g., without updating model parameters using supervised, unsupervised, or semi-supervised learning.

Example Decoder-Based Generative Language Model

[0029] FIG. 1 illustrates an exemplary generative language model **100** (e.g., a transformer-based decoder) that can be employed using the disclosed implementations. Generative language model **100** is an example of a machine learning model that can be used to perform one or more natural language processing tasks that involve generating text, as discussed more below. For the purposes of this document, the term “natural language” means language that is normally used by human beings for writing or conversation.

[0030] Generative language model **100** can receive input text **110**, e.g., a prompt from a user. For instance, the input text can include words, sentences, phrases, or other representations of language. The input text can be broken into tokens and mapped to token and position embeddings **101** representing the input text. Token embeddings can be represented in a vector space where semantically-similar and/or syntactically-similar embeddings are relatively close to one another, and less semantically-similar or less syntactically-similar tokens are relatively further apart. Position embeddings represent the location of each token in order relative to the other tokens from the input text.

[0031] The token and position embeddings **101** are processed in one or more decoder blocks **112**. Each decoder block implements masked multi-head self-attention **103**, which is a mechanism relating different positions of tokens within the input text to compute the similarities between those tokens. Each token embedding is represented as a weighted sum of other tokens in the input text. Attention is only applied for already-decoded values, and future values are masked. Layer normalization **104** normalizes features to mean values of 0 and variance to 1, resulting in smooth

gradients. Feed forward layer **105** transforms these features into a representation suitable for the next iteration of decoding, after which another layer normalization **106** is applied. Multiple instances of decoder blocks can operate sequentially on input text, with each subsequent decoder block operating on the output of a preceding decoder block. After the final decoding block, text prediction layer **107** can predict the next word in the sequence, which is output as output text **120** in response to the input text **110** and also fed back into the language model. The output text can be a newly-generated response to the prompt provided as input text to the generative language model. [0032] Generative language model **100** can be trained using techniques such as next-token prediction or masked language modeling on a large, diverse corpus of documents. For instance, the text prediction layer **107** can predict the next token in a given document, and parameters of the decoder block **112** and/or text prediction layer can be adjusted when the predicted token is incorrect. In some cases, a generative language model can be pretrained on a large corpus of documents and then tuned to a particular use case. For instance, a pretrained generative language model can be tuned using a reinforcement learning technique such as reinforcement learning from human feedback (“RLHF”).

Reinforcement Learning Agents

[0033] In reinforcement learning, an agent can determine a probability distribution over one or more actions that can be taken within an environment, and/or select a specific action to take. An agent can determine the probability distribution and/or select the actions according to a policy. For instance, the policy can map state to probabilities for actions that can be taken by the agent. The agent can refine the policy using a reinforcement learning model that updates the policy based on feedback to actions selected by the agent.

[0034] A reinforcement learning model can be trained using an algorithm to learn a policy using a reward function. The reinforcement learning model can update learnable parameters by receiving feedback and evaluating the reactions using the reward function. For instance, reinforcement learning policies can be implemented using weights that can be learned by training a machine learning model, such as a linear model or neural network.

[0035] A reinforcement learning model can also have hyperparameters that control how the agent acts and/or learns. For instance, a reinforcement learning model can have a learning rate, an exploration strategy, an update frequency, etc. A policy is a function used to determine what actions that an agent takes in a given context. A policy can be learned using reinforcement learning according to a reward function. An agent can utilize state in order to choose which action to take. For instance, a contextual bandit receives context features describing the current state of the environment and uses these features to select the next action to take. A contextual bandit agent can keep a history of rewards earned for different actions taken in different states and continue to modify the policy as new information is discovered.

[0036] One type of contextual bandit is a linear model, such as Vowpal Wabbit. Such a model may output, at each step, a probability density function over the available actions, and select an action randomly from the probability density function. The model may learn feature weights that are applied to one or more input features (e.g., describing context) to determine the probability density function. When the reward obtained in a given step does not match the expected reward, the agent can update the weights used to determine the probability density function.

Example Reinforcement Learning Framework

[0037] FIG. 2 shows an example where an agent **202** receives state information **204**, action information **206**, and feedback information **208**. The state information include a history of previous or current inputs, a state of an application environment **210**, etc. The action information represents one or more available actions **212**. The agent can choose a selected action **214** based on the state information. The feedback information can represent how the state of the environment changes in response to the action selected by the agent or feedback from other sources. The feedback information **208** can be used in a reward function to determine a reward for the agent **202** for each

selected action.

[0038] In some cases, the actions available to an agent can be independent of the state—e.g., all actions can be available to the agent in all states. In other cases, the actions available to an agent can be constrained by the current state, so that actions available to the agent in one state are not available in another state. Thus, in some implementations, state information **204** can specify what the available actions are for an agent given the current state in which the agent is operating.

Example Refinement Workflow

[0039] FIG. **3** shows an example refinement workflow **300**. In FIG. **3**, dotted lines represent online or offline training processing, and solid lines represent inference processing. This section primarily describes inference processing, with further discussion of offline and online training processing provided in sections further below.

[0040] As shown in FIG. **3**, a first machine learning model **302** interacts with an environment **304**. At a time t , the environment has a state $S_{sub.t}$ which is provided to the first machine learning model. The first machine learning model generates a first action $A_{sub.1t}$. A refinement model **306** receives the state and the first action and produces a refined action $A_{sub.2t}$. The first action and the refined action are evaluated by a reward model **308**. An expected reward $R_{sub.1}(t+1)$ is determined for the first action and an expected reward $R_{sub.2}(t+1)$ is determined for the refined action. The expected rewards are output to a policy **310**, which selects a selected action $A_{sub.t}$ and then the environment state is updated to $S_{sub.t+1}$.

[0041] In some cases, a sentiment model **312** can provide sentiment that can be employed to train the reward model. For instance, the sentiment model can be a transformer-based model that evaluates user feedback (e.g., text) to determine whether the feedback has a positive, neutral, or negative sentiment. Also, note that in some cases the environment can provide reward signals that can be employed to train the reward model.

Example Reward and Refinement Models

[0042] FIG. **4A** shows an example of reward model **308**. $R(S_t, A_t)$ is a scalar value showing how well aligned a proposed action $A_{sub.t}$ is to state $S_{sub.t}$. The value is determined using deep neural network weights **402**. The better aligned the proposed action is to the state, the higher the reward value.

[0043] FIG. **4B** shows an example of refinement model **306**. Given a state $S_{sub.t}$ and an action $A_{sub.1t}$, the refinement agent outputs refined action $A_{sub.2t}$. The refined action can be determined using deep neural network weights **404**.

[0044] In some implementations, deep neural network weights **402** and **404** can be completely separate, i.e., the reward model **308** and the refinement model **306** can be completely separate models. In other implementations, the reward model and the refinement model can share certain parameters. For instance, the refinement model can be implemented as a decoder-based generative language model. The reward model can receive the embedding of the last token output by the refinement model and pass that token through a fully connected layer to generate a real valued score reflecting the predicted reward for the refined action. The last token has sufficient information for the reward prediction because the last token has attended to all previous tokens, which makes its embedding fully context aware.

Specific Implementation

[0045] Refinement workflow **300** is a general process that can be employed for a wide range of use cases. The following provides a specific example implementation of refinement workflow **300**, where the first machine learning model **302** is a generative language model with fixed parameters, e.g., a GPT-based model, that provides a Copilot for a user to query a backend service. Here, the environment **304** is a user interface that allows a user to interact with the first machine learning model, the refinement model **306**, and/or the backend service. The refinement model is another generative language model with tunable parameters, such as a LLaMA model. Note that the term “fixed” refers to parameters that are not readily available for tuning, e.g., for a model such as a

version of GPT that is accessible for inference but not generally accessible for training by end users. In contrast, the term “tunable” refers to parameters that can readily be tuned by an end user. For instance, various versions of LLaMA can readily be obtained and tuned by an end user. [0046] In this specific example, the state is one or more queries input by the user. The action space is a set of available application programming interface calls to a pharmaceutical database (the backend service). The reward for a given action can be determined by the reward model **308**. For example, as discussed more below, the reward model can be trained using training signals such as sentiment from sentiment model **312**, how long sessions last, and/or offline training data as discussed more below. The policy **310** can balance exploitation and exploration by, in some cases, using an epsilon-greedy strategy to select from the first actions and refined actions. In some cases, the policy selects the action with the highest reward (exploitation) and in other cases, selects actions with lower reward values to expand the knowledge of the reward model and/or refinement model.

[0047] The reward model **308** can be employed to train the refinement model **306**. Note that the reward model can receive text representations of queries, context, and/or actions generated by the first machine learning model **302** and/or refinement model **306**. Alternatively, the reward model can receive embeddings representing the queries, context, and/or actions generated by the first machine learning model and/or refinement model. The reward model can use either the text representations or the embeddings to predict rewards for actions provided by the first model or the refinement model.

Example Prompt

[0048] FIG. 5 shows an example communication **500**, which represents how communication can be performed with the first machine learning model **302**. The example communication includes a system message **502**, which conveys certain context for the first machine learning model to employ in generating a response. Here, the system message conveys a role for the first machine learning model (an AI assistant), a list of application programming interfaces available to the first machine learning model (shown in FIG. 5 as a placeholder in braces “{ }”), and some example solutions of how the available application programming interfaces can be called to generate solutions to user queries. In addition, placeholders for a current user query **504** and output **506** of the first machine learning model (e.g., a first action representing a solution to the current user query) are also shown. Note that “bkg” in FIG. 5 refers to “biomedical knowledge graph,” a graph structure representing biological and medical knowledge relating to diseases, disease treatments, etc.

Example User Experience

[0049] FIG. 6 shows an example user experience via a graphical user interface **600**. A user inputs a query **602** into the graphical user interface. The query requests a list of genes associated with Type 1 diabetes and drugs that can treat Type 1 diabetes. When the user enters the query, inference processing from refinement workflow **300** executes in the background, resulting in a selected action **604** which is output in response to the query. In this case the selected action is a three-step solution with each step invoking individual application programming interfaces from those that were identified via the prompts described previously. Note that in this example, the user is not aware of whether the selected action was generated by the first machine learning model or the refinement model.

[0050] If the user wants to execute the selected action **604**, the user can click submit button **606**. This will cause the individual application programming interfaces to be executed by the application. Then, the user can be taken to a results screen (not shown) with the results of their query. If the user prefers to see an alternative action proposed, the user can select regenerate button **608**, which can trigger inference processing from refinement workflow **300** to execute again and generate a new selected action. Note that whether the user chooses to submit the solution to the backend service can be used as a training signal, e.g., with submitting the solution being a positive training example for the refinement and/or reward model and regenerating the solution being a

negative training signal for the refinement and/or reward model.

Offline Training of Reward Model

[0051] Referring back to FIG. 4A, the reward model **308** can include weights DNN(w) **402**, which can be learned online or offline. For offline learning, batch learning on a training dataset of previous user interactions with the first machine learning model **302**, the refinement model **306**, or another generative language model can be employed. For instance, a positive reward value can be assigned to previous interactions where a user accepts a selected action output by a generative language model. A negative or zero reward value can be assigned to previous interactions where a user did not accept a selected action output by a generative language model.

[0052] As another example, annotations can be collected from either humans or a generative language model using several approaches. A first approach for collecting annotations can involve:

[0053] For each input state (e.g., a human query), generate two sets of plausible actions (e.g., using the first machine learning model, the refinement model, or another generative language model).

[0054] Send the pairs (S, A.sub.1, A.sub.2) for judgment by a human or generative language model annotator. [0055] The annotator can pick which action is more suitable given the provided input state S.sub.t.

[0056] A second approach for collecting annotations can involve: [0057] For each input state, use the first machine learning model, the refinement model, or another generative language to generate K results. [0058] Send all K results (S, A.sub.1, . . . , A.sub.k) for judgement by a human or generative language model annotator. [0059] The annotator will pick the action that is most appropriate give the state. [0060] This will create K-1, pairwise training instances. Assume i'th action A.sub.i is selected, the pairwise training instances include (S,A.sub.1,A.sub.i), (S, A.sub.2,A.sub.i), . . . , (S, A.sub.K, A.sub.i).

[0061] If the action preferred by the annotator is A.sub.i then a loss function can be defined as:

$$[00001] \text{loss}(r(w)) = -E_{(S, A_0, A_1, i)} [\log (r_w(S, A_i) - r_w(S, A_{1-i}))]$$

which can be computed over the training data. Here, r(w) refers to the weights of the reward model, E refers to an energy function, and σ refers to the sigmoid function, which is taken over a difference between the reward for the action selected by the annotator and the reward for the other action in a given pair of training instances.

[0062] Offline training data can be collected via an application that interfaces with an annotator to select the appropriate action. For a human annotator, the application can include a graphical user interface that displays the proposed actions to a user and receives a user selection (e.g., via a mouse or touch input) of the selected action. For a generative language model annotator, the application can use application programming interfaces associated with the generative language model to provide the proposed solutions and associated context so that the generative language model can select one of the proposed actions.

Online Training of Reward Model

[0063] Other implementations can learn weights DNN(w) **402** of reward model **308** online. For instance, users can provide explicit approval (e.g., thumbs-up or deciding to submit a proposed action) or disapproval (e.g., thumbs-down or deciding not to submit a proposed action) indications. As another example, user messages can be analyzed using sentiment model **312** to infer user sentiment with respect to proposed actions. For these training examples, if the action preferred by a human is A.sub.i then a loss function can be defined as:

$$[00002] \text{loss}(r(w)) = -E_{(S, A_0, A_1, i)} [\log (r_w(S_1, A_1) - r_w(S_0, A_0))]$$

For instance, the loss function and model updates to the reward model weights could happen at the end of each user session, or it could be scheduled to happen periodically (e.g., on a daily or weekly basis).

Offline Training of Refinement Model

[0064] Referring back to FIG. 4B, the refinement model **306** can include weights DNN(w) **404**, which can be pre-trained offline and finetuned online. For offline training of the refinement model,

the following approach can be employed:

[0065] Offline Training Data Collection: Use a prompt library that includes N pairs of queries and responses (Q, R) [0066] For each pair ask the first machine learning model or another generative machine learning model to rewrite Q in M versions [0067] If there is a noun (e.g., gene, drug, chemical, etc. for a pharmaceutical application), replace the noun with K other corresponding values respectively (i.e., a gene, drug, or chemical). [0068] Input each created query to the first machine learning model or the other generative language model and get response R [0069] Then input the query along with the response $R_{sub.1}$ to the current refinement model to get refined response $R_{sub.2}$

This approach results in $N \times M \times K$ training examples for training the refinement model offline. For instance, the refinement model can be trained using a contextual bandit policy optimization loss function with reward values calculated by the reward model for $R_{sub.1}$ and $R_{sub.2}$.

Refinement Model—Online Training

[0070] Other implementations can learn weights $DNN(w)$ **404** of refinement model **308** online. Continuous training for each user interaction can be employed using a reward score provided by the reward model **308** for each action output by the refinement model. For instance, episodic policy optimization, averaged-reward policy optimization, and/or discounted-average-reward policy optimization can be employed to update the refinement model.

[0071] When updating the refinement model **308**, several approaches can be employed to increase robustness. For instance, trust region policy optimization can be employed, where the difference between subsequent policies generated by the refinement model can be constrained to increase stability of optimization. Another approach to increasing robustness involves a proximal policy optimization approach.

Policy

[0072] Generally, policy **310** can choose between actions generated by the first machine learning model and the refinement model according to one or more criteria. In some implementations, policy **310** employs a dynamic epsilon-greedy type method to handle exploration/exploitation. First, the output of the first machine learning model ($A_{sub.1}$) can be selected with probability $1-p$. With prob p consider choosing refinement model output $A_{sub.2}$ as follows:

[0073] If $R(A_{sub.2}, S) \geq R(A_{sub.1}, S)$ [0074] pick $A_{sub.2}$ with prob proportional to $\exp(T^*(R_{sub.2} - R_{sub.1}))$

[0075] Else backoff to use $A_{sub.1}$

Thus, the greater the reward value for $A_{sub.2}$, the more likely $A_{sub.2}$ is to be selected as the selected action.

Example System

[0076] The present implementations can be performed in various scenarios on various devices. FIG. 7 shows an example system **700** in which the present implementations can be employed, as discussed more below.

[0077] As shown in FIG. 7, system **700** includes a client device **710**, a server **720**, a server **730**, and a server **740**, connected by one or more network(s) **750**. Note that the client device can be embodied as a mobile device such as a smart phone or tablet, as well as a stationary device such as a desktop, server device, etc. Likewise, the servers can be implemented using various types of computing devices. In some cases, any of the devices shown in FIG. 7, but particularly the servers, can be implemented in data centers, server farms, etc.

[0078] Certain components of the devices shown in FIG. 7 may be referred to herein by parenthetical reference numbers. For the purposes of the following description, the parenthetical (1) indicates an occurrence of a given component on client device **710**, (2) indicates an occurrence of a given component on server **720**, (3) indicates an occurrence on server **730**, and (4) indicates an occurrence on server **740**. Unless identifying a specific instance of a given component, this document will refer generally to the components without the parenthetical.

[0079] Generally, the devices **710**, **720**, **730**, and/or **740** may have respective processing resources **701** and storage resources **702**, which are discussed in more detail below. The devices may also have various modules that function using the processing and storage resources to perform the techniques discussed herein. The storage resources can include both persistent storage resources, such as magnetic or solid-state drives, and volatile storage, such as one or more random-access memory devices. In some cases, the modules are provided as executable instructions that are stored on persistent storage devices, loaded into the random-access memory devices, and read from the random-access memory by the processing resources for execution.

[0080] Client device **710** can include a local assistant application **711** that can interact with a remote assistant application **721** on server **720**. For instance, the local assistant application can display a graphical user interface such as graphical user interface **600**, shown in FIG. **6**. The local assistant application can send received queries to the remote assistant application. The remote assistant application can provide the received queries to the first machine learning model **302** on server **730** and receive a first action from the first machine learning model.

[0081] The remote assistant application can also include the refinement model **306**, the reward model **308**, the policy **310**, and/or the sentiment model **312** which can be employed for refinement workflow **300** as described previously. The refinement model can output a refined action, and the policy can choose either the first action or the refined action. Then, the selected action can be sent to client device **710**. If the user requests to invoke the selected action, the selected action can be invoked by calling one or more application programming interfaces on service **741**. A reinforcement learning updater **722** on server **720** can update the reward model **308** and/or refinement model **306**, e.g., using any of the techniques described previously.

Example Method

[0082] FIG. **8** illustrates an example method **800**, consistent with the present concepts. As discussed more below, method **800** can be implemented on many different types of devices, e.g., by one or more cloud servers, by a client device such as a laptop, tablet, or smartphone, or by combinations of one or more servers, client devices, etc.

[0083] Method **800** begins at block **802**, where a state of an environment is received. For instance, as noted above, the state can be a query (e.g., a natural language query) relating to a service. In some cases, the state can include other information, such as previous queries in the same session, previous actions that have been accepted or rejected during the session, etc. As another example, the state can include context such as a user profile, location, time, etc.

[0084] Method **800** continues at block **804**, where the state is input to a first machine learning model. For instance, the first machine learning model can be a generative language model with fixed parameters. As noted previously, GPT is one example of a generative language model that may have parameters that are not readily tuned by end users.

[0085] Method **800** continues at block **806**, where a first action is received from the first machine learning model. For instance, the first action can be a proposed solution of one or more application programming interfaces to invoke on a service to satisfy the query. The first action can be generated based on the query as well as any other state information as described above with respect to block **802**.

[0086] Method **800** continues at block **808**, where the state and the first action are input to a refinement model. For instance, the refinement model can be a generative language model with tunable parameters. As noted previously, LLaMA is one example of a generative language model that may have parameters that are accessible for tuning by end users.

[0087] Method **800** continues at block **810**, where a refined action is received from the refinement model. For instance, the refined action can be another solution of one or more application programming interfaces to invoke on a service to satisfy the query. The refined action can be generated based on the query, the first action, as well as any other state information as described above with respect to block **802**.

[0088] Method **800** continues at block **812**, where the state, the first action, and the refined action are input to a reward model. For instance, the reward model can be a deep neural network that can, in some cases, share some parameters (e.g., weights) with the refinement model.

[0089] Method **800** continues at block **814**, where reward values are received for the first action and the refined action. For instance, as noted above, the reward values can reflect how well the first action and the refined action are with the state. Higher reward values can convey that a given action is well-aligned with the state, and lower reward values can convey that a given action is not well-aligned with the state.

[0090] Method **800** continues at block **816**, where an action is selected from the first action and the refined action. For instance, as noted above, the action can be selected based on a policy that balances exploitation vs. exploration. In some cases, the policy can preferentially (although not necessarily deterministically) select the action with a higher reward value.

[0091] Method **800** continues at block **818**, where the selected action is output. For instance, the selected action can be sent to a user that submitted the query, where the user has the option to invoke the selected action or else request another action. In some cases, block **818** can also involve updating the reward model and/or refinement model based at least on whether the user accepts the selected action.

Additional Implementations

[0092] The concepts described herein have been explained above with respect to a scenario where states relate to user queries and the actions relate to application interface calls output by generative language models to satisfy the user queries. However, the concepts described herein can be employed in a wide range of application scenarios beyond those described above. The following describes a few examples of how the present concepts can be implemented in other application scenarios.

[0093] Consider a strategy game scenario (e.g., chess) where a user is playing against a computer opponent. The computer opponent could be implemented using refinement workflow **300** as described above, with two different machine learning models receiving game state (e.g., locations of chess pieces) and selecting actions to take based on the game state. One of the machine learning models could have fixed parameters and output a first action (e.g., moving a chess piece from one square on the board to another). The second machine learning model could be a tunable refinement model and output a refined action (e.g., a different chess move) based on the first action and the game state. A policy could be employed to select which action to take, where the reward values for each action could reflect information such as whether an opponent's piece is taken by a proposed action, the rank of any pieces that are taken, whether the move puts the opponent in check or checkmate, etc.

[0094] As another example, consider an application that assists a software developer in writing code. An integrated development environment could implement refinement workflow **300** as described above, with two different machine learning models receiving a natural language query requesting that a function be written in a programming language to calculate a particular numeric value. One of the machine learning models could have fixed parameters and output a first action (e.g., a recursive function). The second machine learning model could be a tunable refinement model and output a refined action (e.g., a function with a while loop) based on the first action and the natural language query. A policy could be employed to select which action to take, where the reward values for each action could reflect information such as whether the function compiles properly, the amount of stack memory used by the function, etc.

TECHNICAL EFFECT

[0095] As noted previously, it is possible to train or tune generative language models or other machine learning models for specific tasks. Moreover, tuning a pretrained model to a new task domain can be a very efficient approach compared to training a new model from scratch, particularly for large models such as generative language models that are pretrained on massive

amounts of data. However, in some cases it is not technically feasible to train or tune a machine learning model for a specific task, e.g., because the model is only made available to end users for inference processing and the parameters are not accessible to the end users for tuning.

[0096] By using a refinement model to refine the actions output by a first machine learning model with fixed parameters as disclosed herein, the refinement model can learn over time to improve upon actions output by the first machine learning model. Furthermore, by training a reward model to predict expected rewards for actions output by both models, the disclosed techniques can adaptively select between actions output by both models.

Device Implementations

[0097] As noted above with respect to FIG. 7, system **700** includes several devices, including a client device **710**, a server **720**, a server **730**, and a server **740**. As also noted, not all device implementations can be illustrated, and other device implementations should be apparent to the skilled artisan from the description above and below.

[0098] The term “device,” “computer,” “computing device,” “client device,” and or “server device” as used herein can mean any type of device that has some amount of hardware processing capability and/or hardware storage/memory capability. Processing capability can be provided by one or more hardware processors (e.g., hardware processing units/cores) that can execute data in the form of computer-readable instructions to provide functionality. Computer-readable instructions and/or data can be stored on storage, such as storage/memory and or the datastore. The term “system” as used herein can refer to a single device, multiple devices, etc.

[0099] Storage resources can be internal or external to the respective devices with which they are associated. The storage resources can include any one or more of volatile or non-volatile memory, hard drives, flash storage devices, and/or optical storage devices (e.g., CDs, DVDs, etc.), among others. As used herein, the term “computer-readable media” can include signals. In contrast, the term “computer-readable storage media” excludes signals. Computer-readable storage media includes “computer-readable storage devices.” Examples of computer-readable storage devices include volatile storage media, such as RAM, and non-volatile storage media, such as hard drives, optical discs, and flash memory, among others.

[0100] In some cases, the devices are configured with a general purpose hardware processor and storage resources. Processors and storage can be implemented as separate components or integrated together as in computational RAM. In other cases, a device can include a system on a chip (SOC) type design. In SOC design implementations, functionality provided by the device can be integrated on a single SOC or multiple coupled SOC. One or more associated processors can be configured to coordinate with shared resources, such as memory, storage, etc., and/or one or more dedicated resources, such as hardware blocks configured to perform certain specific functionality. Thus, the term “processor,” “hardware processor” or “hardware processing unit” as used herein can also refer to central processing units (CPUs), graphical processing units (GPUs), controllers, microcontrollers, processor cores, or other types of processing devices suitable for implementation both in conventional computing architectures as well as SOC designs.

[0101] Alternatively, or in addition, the functionality described herein can be performed, at least in part, by one or more hardware logic components. For example, and without limitation, illustrative types of hardware logic components that can be used include Field-programmable Gate Arrays (FPGAs), Application-specific Integrated Circuits (ASICs), Application-specific Standard Products (ASSPs), System-on-a-chip systems (SOCs), Complex Programmable Logic Devices (CPLDs), etc.

[0102] In some configurations, any of the modules/code discussed herein can be implemented in software, hardware, and/or firmware. In any case, the modules/code can be provided during manufacture of the device or by an intermediary that prepares the device for sale to the end user. In other instances, the end user may install these modules/code later, such as by downloading executable code and installing the executable code on the corresponding device.

[0103] Also note that devices generally can have input and/or output functionality. For example,

computing devices can have various input mechanisms such as keyboards, mice, touchpads, voice recognition, gesture recognition (e.g., using depth cameras such as stereoscopic or time-of-flight camera systems, infrared camera systems, RGB camera systems or using accelerometers/gyroscopes, facial recognition, etc.). Devices can also have various output mechanisms such as printers, monitors, etc.

[0104] Also note that the devices described herein can function in a stand-alone or cooperative manner to implement the described techniques. For example, the methods and functionality described herein can be performed on a single computing device and/or distributed across multiple computing devices that communicate over network(s) 750. Without limitation, network(s) 750 can include one or more local area networks (LANs), wide area networks (WANs), the Internet, and the like.

[0105] Various examples are described above. Additional examples are described below. One example includes a computer-implemented method comprising obtaining a state of an environment, inputting the state to a first machine learning model, receiving a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state, inputting the state and the first action to a refinement model, receiving a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action, providing the state, the first action, and the refined action to a reward model, receiving, from the reward model, respective reward values for the first action and the refined action, based at least on the reward values, selecting the first action or the refined action as a selected action, and outputting the selected action.

[0106] Another example can include any of the above and/or below examples where the first machine learning model comprises a first generative language model.

[0107] Another example can include any of the above and/or below examples where the refinement model comprises a second generative language model.

[0108] Another example can include any of the above and/or below examples where the state includes a user query.

[0109] Another example can include any of the above and/or below examples where the first action and the refined action include one or more application programming interfaces provided by a service.

[0110] Another example can include any of the above and/or below examples where the method further comprises inputting a first prompt to the first generative language model instructing the first generative language model to generate the first action based at least on the user query, available application programming interfaces, and one or more example solutions to example user queries, the example solutions utilizing the available application programming interfaces.

[0111] Another example can include any of the above and/or below examples where the first generative language model and the second generative language models being decoder-based models.

[0112] Another example can include any of the above and/or below examples where the method further comprises performing offline training of the reward model by generating multiple actions for a given state, receiving annotations from a human or machine learning model for the multiple actions, and training the reward model based at least on the annotations.

[0113] Another example can include any of the above and/or below examples where the method further comprises performing online training of the reward model based at least on explicit feedback or sentiment analysis using a sentiment model.

[0114] Another example can include any of the above and/or below examples where the method further comprises performing offline training of the refinement model based at least on reward values from the reward model for refined actions generated by the refinement model when refining actions present in a training data set.

[0115] Another example can include any of the above and/or below examples where the method

further comprises performing online training of the refinement model based at least on reward values from the reward model for refined actions generated by the refinement model when interacting with a user.

[0116] Another example can include any of the above and/or below examples where the refinement model and the reward model share one or more parameters.

[0117] Another example can include any of the above and/or below examples where the reward model computes rewards over particular embeddings output by the refinement model.

[0118] Another example can include any of the above and/or below examples where the particular embeddings being embeddings of final tokens output by the refinement model.

[0119] Another example can include any of the above and/or below examples where the first machine learning model having fixed parameters.

[0120] Another example can include a system comprising a hardware processing unit, and a storage resource storing computer-readable instructions which, when executed by the hardware processing unit, cause the system to obtain a state of an environment, input the state to a first machine learning model, receive a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state, input the state and the first action to a refinement model, receive a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action, provide the state, the first action, and the refined action to a reward model, receive, from the reward model, respective reward values for the first action and the refined action, based at least on the reward values, select the first action or the refined action as a selected action, and output the selected action.

[0121] Another example can include any of the above and/or below examples where the computer-readable instructions, when executed by the hardware processing unit, cause the system to train the reward model using training data generated with a generative language model.

[0122] Another example can include any of the above and/or below examples where the computer-readable instructions, when executed by the hardware processing unit, cause the system to train the refinement model using rewards generated by the reward model.

[0123] Another example can include any of the above and/or below examples where the computer-readable instructions, when executed by the hardware processing unit, cause the system to extract one or more application programming interfaces from the selected action, and invoke the one or more application programming interfaces on a service.

[0124] Another example can include computer-readable storage medium storing computer-readable instructions which, when executed by a processing unit, cause the processing unit to perform acts comprising obtaining a state of an environment, inputting the state to a first machine learning model, receiving a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state, inputting the state and the first action to a refinement model, receiving a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action, providing the state, the first action, and the refined action to a reward model, receiving, from the reward model, respective reward values for the first action and the refined action, based at least on the reward values, selecting the first action or the refined action as a selected action, and outputting the selected action.

CONCLUSION

[0125] Although the subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as example forms of implementing the claims and other features and acts that would be recognized by one skilled in the art are intended to be within the scope of the claims.

Claims

1. A computer-implemented method comprising: obtaining a state of an environment; inputting the state to a first machine learning model; receiving a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state; inputting the state and the first action to a refinement model; receiving a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action; providing the state, the first action, and the refined action to a reward model; receiving, from the reward model, respective reward values for the first action and the refined action; based at least on the reward values, selecting the first action or the refined action as a selected action; and outputting the selected action.
2. The method of claim 1, wherein the first machine learning model comprises a first generative language model.
3. The method of claim 2, wherein the refinement model comprises a second generative language model.
4. The method of claim 3, wherein the state includes a user query.
5. The method of claim 4, wherein the first action and the refined action include one or more application programming interfaces provided by a service.
6. The method of claim 5, further comprising: inputting a first prompt to the first generative language model instructing the first generative language model to generate the first action based at least on the user query, available application programming interfaces, and one or more example solutions to example user queries, the example solutions utilizing the available application programming interfaces.
7. The method of claim 6, the first generative language model and the second generative language models being decoder-based models.
8. The method of claim 1, further comprising: performing offline training of the reward model by generating multiple actions for a given state, receiving annotations from a human or machine learning model for the multiple actions, and training the reward model based at least on the annotations.
9. The method of claim 1, further comprising: performing online training of the reward model based at least on explicit feedback or sentiment analysis using a sentiment model.
10. The method of claim 1, further comprising: performing offline training of the refinement model based at least on reward values from the reward model for refined actions generated by the refinement model when refining actions present in a training data set.
11. The method of claim 1, further comprising: performing online training of the refinement model based at least on reward values from the reward model for refined actions generated by the refinement model when interacting with a user.
12. The method of claim 1, wherein the refinement model and the reward model share one or more parameters.
13. The method of claim 12, wherein the reward model computes rewards over particular embeddings output by the refinement model.
14. The method of claim 13, the particular embeddings being embeddings of final tokens output by the refinement model.
15. The method of claim 1, the first machine learning model having fixed parameters.
16. A system comprising: a hardware processing unit; and a storage resource storing computer-readable instructions which, when executed by the hardware processing unit, cause the system to: obtain a state of an environment; input the state to a first machine learning model; receive a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state; input the state and the first action to a refinement model; receive a

refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action; provide the state, the first action, and the refined action to a reward model; receive, from the reward model, respective reward values for the first action and the refined action; based at least on the reward values, select the first action or the refined action as a selected action; and output the selected action.

17. The system of claim 16, wherein the computer-readable instructions, when executed by the hardware processing unit, cause the system to: train the reward model using training data generated with a generative language model.

18. The system of claim 17, wherein the computer-readable instructions, when executed by the hardware processing unit, cause the system to: train the refinement model using rewards generated by the reward model.

19. The system of claim 16, wherein the computer-readable instructions, when executed by the hardware processing unit, cause the system to: extract one or more application programming interfaces from the selected action; and invoke the one or more application programming interfaces on a service.

20. A computer-readable storage medium storing computer-readable instructions which, when executed by a processing unit, cause the processing unit to perform acts comprising: obtaining a state of an environment; inputting the state to a first machine learning model; receiving a first action from the first machine learning model, the first machine learning model selecting the first action based at least on the state; inputting the state and the first action to a refinement model; receiving a refined action from the refinement model, the refinement model selecting the refined action based at least on the state and the first action; providing the state, the first action, and the refined action to a reward model; receiving, from the reward model, respective reward values for the first action and the refined action; based at least on the reward values, selecting the first action or the refined action as a selected action; and outputting the selected action.
