



US 20250265482A1

(19) **United States**

(12) **Patent Application Publication**
DOTAN-COHEN et al.

(10) **Pub. No.: US 2025/0265482 A1**

(43) **Pub. Date: Aug. 21, 2025**

(54) **ENHANCED COMPUTER EXPERIENCE
FROM ACTIVITY PREDICTION**

G06Q 10/101 (2023.01)

H04W 4/029 (2018.01)

(71) Applicant: **Microsoft Technology Licensing, LLC,**
Redmond, WA (US)

(52) **U.S. Cl.**

CPC *G06N 5/048* (2013.01); *G06F 3/01*
(2013.01); *G06F 3/023* (2013.01); *G06F*
3/038 (2013.01); *G06F 3/0481* (2013.01);
G06F 9/453 (2018.02); *G06N 7/01* (2023.01);
G06N 20/00 (2019.01); *G06Q 10/04*
(2013.01); *G06Q 10/101* (2013.01); *H04W*
4/029 (2018.02); *Y02D 30/70* (2020.08)

(72) Inventors: **Dikla DOTAN-COHEN,** Herzliya (IL);
Shira WEINBERG, Tel Aviv (IL)

(21) Appl. No.: **19/070,296**

(22) Filed: **Mar. 4, 2025**

Related U.S. Application Data

(63) Continuation of application No. 17/345,233, filed on
Jun. 11, 2021, now abandoned.

Publication Classification

(51) **Int. Cl.**

G06N 5/048 (2023.01)

G06F 3/01 (2006.01)

G06F 3/023 (2006.01)

G06F 3/038 (2013.01)

G06F 3/0481 (2022.01)

G06F 9/451 (2018.01)

G06N 7/01 (2023.01)


G06N 20/00 (2019.01)

G06Q 10/04 (2023.01)

(57)

ABSTRACT

A user activity pattern may be ascertained using signal data from a set of computing devices. The activity pattern may be used to infer user intent with regards to a user interaction with a computing device or to predict a likely future action by the user. In one implementation, a set of computing devices is monitored to detect user activities using sensors associated with the computing devices. Activity features associated with the detected user activities are determined and used to identify an activity pattern based on a plurality of user activities having similar features. Examples of user activity patterns may include patterns based on time, location, content, or other context. The inferred user intent or predicted future actions may be used to provide improved user experiences, such as personalization, modifying functionality of user devices, or providing more efficient consumption of bandwidth or power.



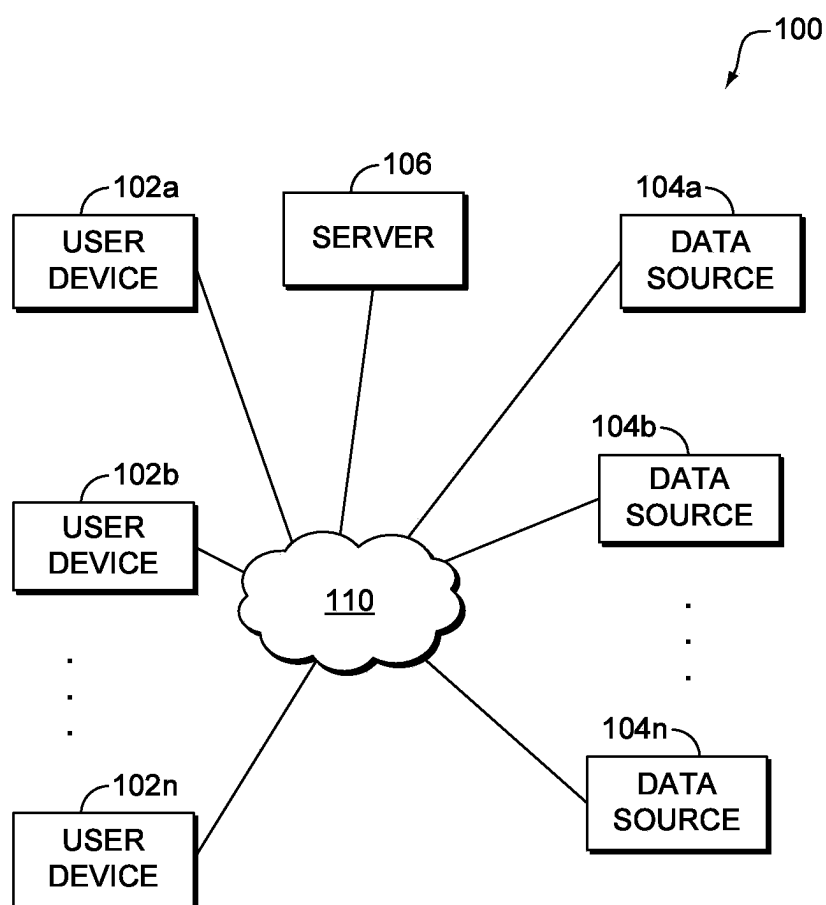


FIG. 1




FIG. 2




FIG. 3

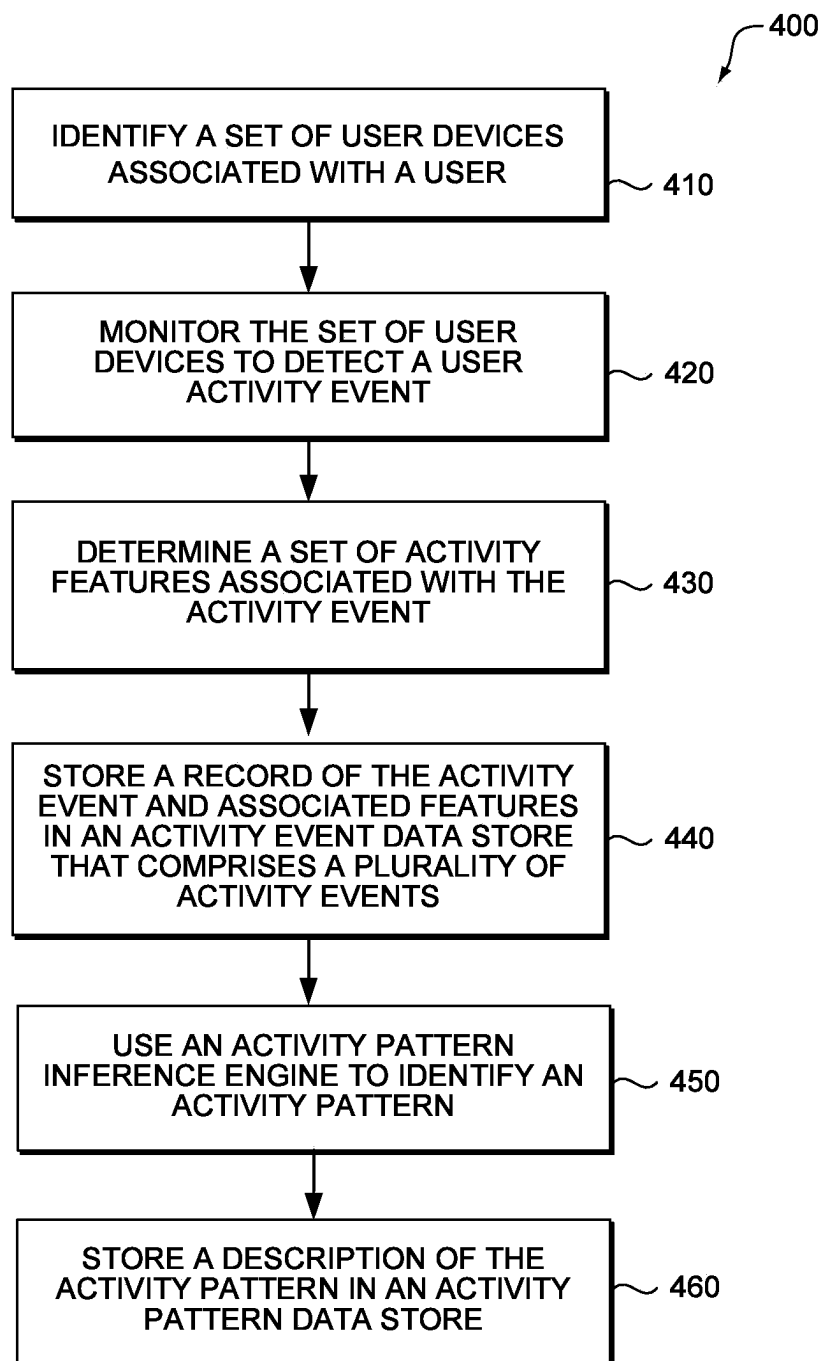


FIG. 4

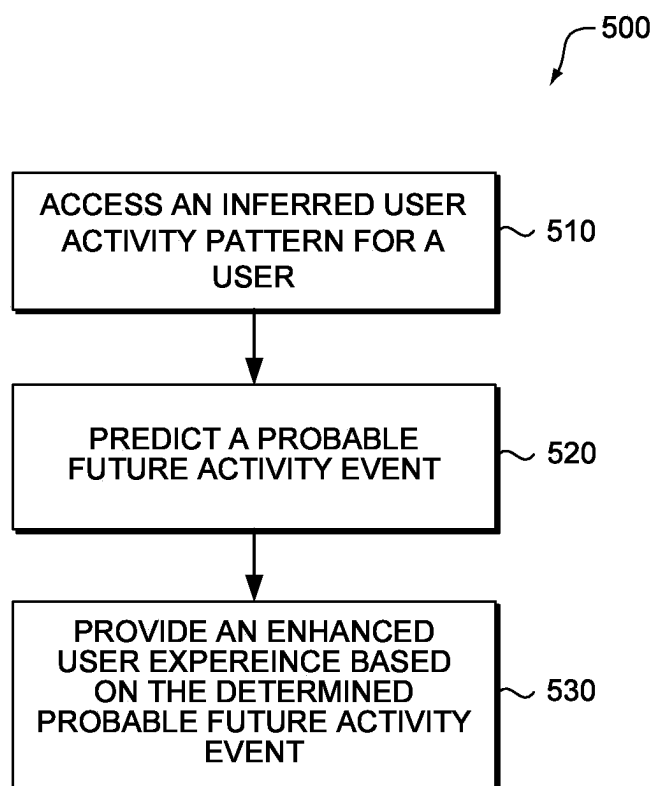


FIG. 5

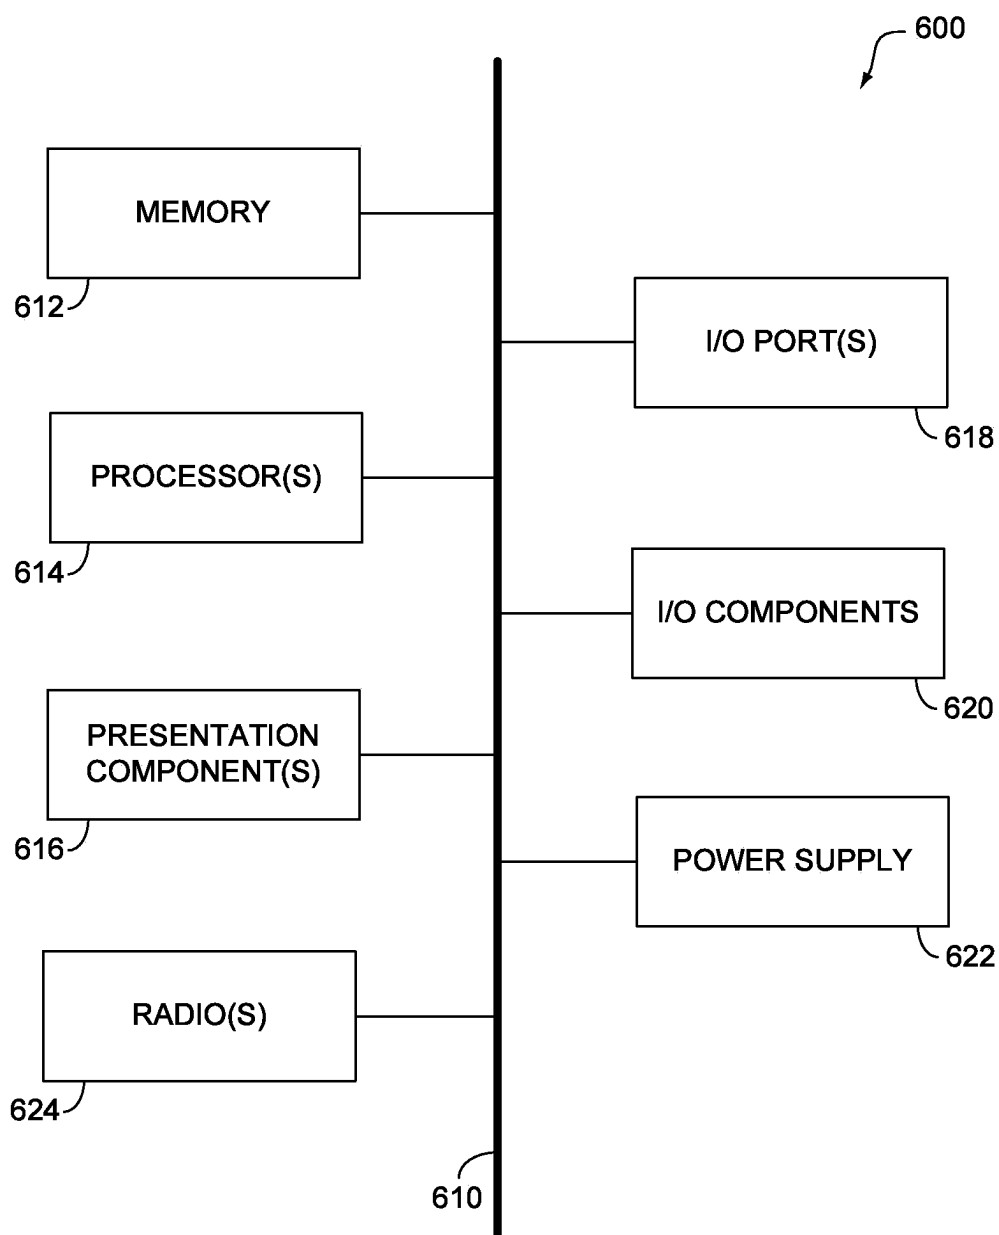


FIG. 6

ENHANCED COMPUTER EXPERIENCE FROM ACTIVITY PREDICTION

CROSS REFERENCE TO RELATED APPLICATIONS

[0001] This application is a continuation of U.S. patent application Ser. No. 17/345,233, filed on Jun. 11, 2021, which itself is a continuation of U.S. Pat. No. 11,429,883 filed on Nov. 13, 2015, which is expressly incorporated herein by reference in its entirety.

BACKGROUND

[0002] People are increasingly interacting with computing devices and relying on these devices for information, recommendations, and other services to assist them in their day-to-day tasks. But understanding a user's intent by the computing device in these interactions remains a difficult technical problem. In such interactions, users are often left frustrated by the inability of their computerized personal assistant applications or services to understand their intent or anticipate their needs.

[0003] At the same time, many users of computing devices have repeating patterns of usage. For example, a user may launch an email app on their mobile device every workday morning, before starting work, browse to a favorite news website over lunchtime, call a close friend or family member on their drive home from work, or use their laptop computer to plan their annual summer vacation around May. By learning these patterns of user activity, the computerized personal assistant applications and services can provide improved user experiences including improved understanding of the user's intent.

SUMMARY

[0004] This summary is provided to introduce a selection of concepts in a simplified form that are further described below in the detailed description. This summary is not intended to identify key features or essential features of the claimed subject matter, nor is it intended to be used in isolation as an aid in determining the scope of the claimed subject matter.

[0005] Embodiments described in the present disclosure are directed towards inferring user intent and predicting future user activity associated with a user's computing devices. In particular, embodiments may determine a likely intent of a user with regards to an interaction with a computing device, or may predict a likely future action by a user, based on a history of sensed user activity. Data corresponding to user activity may be gathered over time using sensors of one or more of the user's computing devices. From this historical user activity information, a computer system may learn user activity patterns associated with the computing devices. By analyzing a user activity pattern, future user actions may be predicted or user intent may be inferred.

[0006] For example, computing devices associated with a user ("user devices") may employ one or more sensors to generate data relevant to a user's activity on a user device(s). The user activity may be monitored, tracked, and used for determining patterns of user activity. Examples of user activity patterns may include, without limitation, activity patterns based on time, location, content, or other context, as described herein. In some embodiments, the activity patterns

may be determined based on user activity related to browsing, application usage, or other related user activities associated with one or more user devices or user activity otherwise determinable via the one or more user devices.

[0007] Based on the determined patterns of user activity, user intent may be inferred regarding user interactions with a computing device, and/or predictions of user activity determined and used to provide improved user experiences. Examples of improved user experiences, which are further described herein, may include personalization services, such as tailoring content for the user; modifying functionality, such as preferences or settings, of user devices associated with the user; or providing more efficient consumption of bandwidth or power of user devices associated with the user. In some embodiments, the user activity patterns, or the inferred intent or predictions of future activity determined therefrom, may be made available to one or more applications and services that consume this information to provide an improved user experience.

[0008] Accordingly, as will be further described, in one embodiment, one or more user devices associated with a user may be identified and monitored for user activity. In some embodiments, user activity monitoring may be facilitated using an application or service that runs on the monitored user device. User actions or "activity events," such as visiting a website, launching an application, or other actions similar to those described herein, may be detected and logged with associated contextual data, for example, by logging the observed user action with a time stamp, location stamp, and/or associating the activity event with other available contextual information. From the logged user activity, historical user activity information may be determined and provided to an inference engine. Based on an analysis of historical user activity, a user activity pattern may be determined. In some embodiments, a semantic analysis is performed on the activity events and associated information to characterize aspects of the activity event. The semantic analysis may be used to characterize information associated with the activity event and may provide other relevant features of the activity events that may be used to identify patterns.

[0009] The user activity patterns may be used to infer user intent or predict future activity of the user. From these predictions or inferred user intent, various implementations may provide personalized computing experiences and other services tailored to the user, such as timely, relevant delivery or presentation of content, or modifying content that would otherwise be presented; improvements to user device performance and network bandwidth usage; incorporation of the user's routine into recommendations and notifications; improved semantic understanding by the user's computer devices; or other examples as described herein. In this way, embodiments of the disclosure provide an improved user experience.

BRIEF DESCRIPTION OF THE DRAWINGS

[0010] Aspects of the disclosure are described in detail below with reference to the attached drawing figures, wherein:

[0011] FIG. 1 is a block diagram of an example operating environment suitable for implementations of the present disclosure;

[0012] FIG. 2 is a diagram depicting an example computing architecture suitable for implementing aspects of the present disclosure;

[0013] FIG. 3 illustratively depicts aspects of an example system for determining user activity patterns based on browser and application activity across multiple user devices, in accordance with an embodiment of the present disclosure;

[0014] FIG. 4 depicts a flow diagram of a method for inferring a user activity pattern, in accordance with an embodiment of the present disclosure;

[0015] FIG. 5 depicts a flow diagram of a method for determining a probable future user action based on a pattern of user activity, in accordance with an embodiment of the present disclosure; and

[0016] FIG. 6 is a block diagram of an exemplary computing environment suitable for use in implementing an embodiment of the present disclosure.

DETAILED DESCRIPTION

[0017] The subject matter of aspects of the present disclosure is described with specificity herein to meet statutory requirements. However, the description itself is not intended to limit the scope of this patent. Rather, the inventors have contemplated that the claimed subject matter might also be embodied in other ways, to include different steps or combinations of steps similar to the ones described in this document, in conjunction with other present or future technologies. Moreover, although the terms “step” and/or “block” may be used herein to connote different elements of methods employed, the terms should not be interpreted as implying any particular order among or between various steps herein disclosed unless and except when the order of individual steps is explicitly described.

[0018] Aspects of the present disclosure relate to inferring user intent and predicting future user activity associated with a user's computing devices. In particular, embodiments described herein may determine a likely intent of a user with regards to an interaction with a user device, or may predict a likely future action by a user, based on a history of sensed user activity. Data corresponding to user activity may be gathered over time using sensors on one or more user devices associated with the user. From this historical user activity information, a computer system may learn user activity patterns associated with the user devices. By analyzing a user activity pattern, future user actions may be predicted or user intent may be inferred. In some cases, the user activity pattern may be analyzed along with sensor data collected by a user device, and the user's intent inferred based on determining a likely intent that is consistent with the determined user activity pattern.

[0019] As further described herein, in some embodiments, user devices may employ one or more sensors to generate data relevant to a user's activity via a user device(s). The user activity may be monitored, tracked, and used for determining user activity patterns. The term “activity pattern” is used broadly herein and may refer to a plurality of user interactions conducted using one or more user devices, activity by the user on or in connection to one or more user devices, events (including actions) related to user activity, or any type of user activity determinable via a computing device, wherein the plurality of interactions, actions, events, or activity share common features or characteristics. In some embodiments, these in-common features or variables may

comprise features characterizing a user activity, time, location, or other contextual information associated with the user activity, as further described herein. Examples of user activity patterns may include, without limitation, activity patterns based on time (e.g., the user browses his bank's website near the beginning of each month to check his account balance), location (e.g., upon arriving at work in the morning, a user turns down the volume on her phone), content (e.g., a user typically browses news-related websites followed by their social-media related websites), or other context, as described herein.

[0020] In some embodiments, the user activity may be related to a user's browsing activity, such as websites, categories of websites, or sequences of websites and/or website categories visited by a user, and user activity associated with the browsing activity. In addition or alternatively, the user activity may be related to a user's application (or app) related activity, such as application usage, which may include usage duration, launches, files accessed via the application or in conjunction with the application usage, or content associated with the application. The term application or app is used broadly herein, and generally refers to a computer program or computer application, which may comprise one or more programs or services, and may run on the user's device(s) or in the cloud.

[0021] Based on the determined patterns of user activity, user intent may be inferred, regarding user interactions with a computing device, or predictions of user activity determined and used to provide improved user experiences. Examples of these improved user experiences, which are further described below, include personalization services, such as tailoring content for the user, more efficient consumption of bandwidth or power, or modifying functionality, such as preferences or settings, of user devices associated with the user. In some embodiments, the user activity patterns, or the inferred intent or predictions of future activity determined therefrom, may be made available to one or more applications and services that consume this information and provide an improved user experience. For example, in one embodiment, the activity pattern information may be provided via an Application Programming Interface (API) so that third-party applications and services can use it, such as by providing timely recommendations, suggestions or other information or services relevant to the user based on the learned activity patterns.

[0022] Accordingly, at a high level, in one embodiment, user data is received from one or more data sources. The user data may be received by collecting user data with one or more sensors or components on user device(s) associated with a user. Examples of user data, also described in connection to component 210 of FIG. 2, may include information about the user device(s), user activity associated with the user devices (e.g., app usage, online activity, searches, calls, usage duration, and other user-interaction data), network-related data (such as network ID, connection data, or other network-related information), application data, contacts data, calendar and social network data, or nearly any other source of user data that may be sensed or determined by a user device or other computing device. The received user data may be monitored and information about the user activity may be stored in a user profile, such as user profile 240 of FIG. 2.

[0023] In some embodiments, based on an identification of one or more user devices, which may be determined from

the user data, the one or more user devices are monitored for user activity. In some embodiments, user activity monitoring may be facilitated using an application or service that runs on the monitored user device. Alternatively or in addition, the user activity monitoring may be facilitated using an application or service that runs in the cloud, which may scan the user device, or detects online activity associated with the user device, such as http requests or other communication information, or otherwise receive information about user activity from the user device.

[0024] User data may be analyzed to detect various features associated with user actions. Detected user actions or “activity events,” which may include actions such as websites visited, applications launched, or other actions similar to those described herein, may be logged with associated contextual data, for example, by logging the observed user action with a corresponding time stamp, location stamp, and/or associating the activity event with other available contextual information. In some embodiments, such logging may be performed on each user device, so that activity patterns may be determined across devices. Further, in some embodiments, cloud-based user-activity information sources may be used such as online user calendars or user activities determined from social media posts, emails, or the like. These sources also may be used for providing other context to the user activity detected on the user device(s). Examples of contextual information are further described in connection to contextual information extractor 284 in FIG. 2. In some embodiments, user activity logs from multiple user devices and available user activity information from cloud-based sources may be combined, thereby representing a composite user activity history. The user activity logs, including corresponding contextual information, may be stored in a user profile associated with the user, such as user profile 240 of FIG. 2.

[0025] From the activity logs or user activity data, historical user activity information may be determined and provided to an inference engine. Based on an analysis of historical user activity, and in some cases current sensor data regarding user activity, a set of one or more likely user activity patterns may be determined. In particular, the inference engine may conduct an analysis of the historical user activity information to identify user activity patterns, which may be determined by detecting repetitions of similar user actions or routines, in some embodiments.

[0026] In some embodiments, a corresponding confidence weight or confidence score may be determined regarding the user activity patterns. The confidence score may be based on the strength of the pattern, which may be determined by the number of observations used to determine a pattern, how frequently the user activity is consistent with the pattern, the age or freshness of the activity observations, the number of features in common with the activity observations that make up the pattern, or similar measurements. In some instances, the confidence score may be considered when providing a personalized user experience or other improved user experience. Further, in some embodiments, a minimum confidence score may be needed before using the activity pattern to provide such experiences or other services. For example, in one embodiment, a threshold of 0.6 (or just over fifty percent) is utilized such that only activity patterns having a 0.6 (or greater) likelihood of predicting user activity may be considered. Nevertheless, where confidence scores and thresholds are used, determined patterns of user activity with

confidence scores less than the threshold may still be monitored since additional observations of user activities may increase the confidence for a particular pattern.

[0027] In some embodiments, crowdsourced user activity history may also be utilized in conjunction with a user’s own activity history. For example, for a given user, a set of other users similar to the given user may be identified, based on having features or characteristics in common with the given user. This might include other users located in proximity to the given user, the given user’s social media friends, work colleagues (which may be determined from an analysis of contextual information associated with the given user), other users with similar user activity patterns, or the like. Information about user activity history from the other users may be relied upon for inferring patterns of user activity for the given user. This may be particularly useful in situations where little user activity history exists for the given user, such as where the user is a new user. In some embodiments, user activity information from similar users may be imputed to the new user until enough user history is available for the new user to determine statistically reliable user pattern predictions, which may be determined based on the number of observations included in the user activity history information or the statistical confidence of the determined user activity patterns, as further described herein. In some cases, where the user activity history comes from other users, the resulting inferred activity patterns for the given user may be assigned a lower confidence.

[0028] In some embodiments, a semantic analysis is performed on the activity events and associated information to characterize aspects of the activity event. For example, features associated with an activity event may be categorized (such as by type, similar time frame or location, for example), and related features may be identified for use in determining a similarity or relational proximity to other activity events, such as by having one or more of the characteristics, including categories and related features, in common. In some embodiments, a semantic knowledge representation, such as a relational knowledge graph, may be employed. In some embodiments, the semantic analysis may use rules, logic such as associations or conditions, or classifiers.

[0029] The semantic analysis may also be used to characterize information associated with the activity event, such as determining that a location associated with the activity corresponds to a hub or venue of interest to the user based on frequency of visits, such as the user’s home, work, gym, etc. (For example, the user’s home hub may be determined to be the location where the user spends most of her time between 8 PM and 6 AM.) Similarly, the semantic analysis may determine time of day that correspond to working hours, lunch time, commute time, etc.

[0030] In this way, the semantic analysis may provide other relevant features of the activity events that may be used to determine patterns. For example, in addition to determining a particular website that the user visited at a certain time, such as visiting CNN.com over lunch, the category of the website may be determined, such as a news-related website. Similarly, the semantic analysis may categorize the activity as being associated with work or home, based on the characteristics of the activity (e.g., a batch of online searches about chi-squared distribution that occurs during working hours at a location corresponding to the user’s office may be determined to be work-related

activity, whereas streaming a movie on Friday night at a location corresponding to the user's home may be determined to be home-related activity). These aspects characterizing the user activity event then may be considered when evaluating the user activity history to identify patterns. For example, a pattern of visiting news-related websites over lunch may be determined where a user routinely visits news-related websites over lunch, but only occasionally visits CNN.com as one of the news-related websites.

[0031] As described previously, the user activity patterns may be used to infer user intent or predict future activity of the user. From these predictions or inferred user intent, various implementations may provide improved user experiences. For example, some embodiments may provide timely, relevant delivery or presentation of content, or modifying content that would otherwise be presented; improvements to user device performance and network bandwidth usage; improved speech recognition; incorporation of the user's routine into recommendations and notifications; improved semantic understanding by the user's computer devices; or other examples of improved user experiences described herein. Some embodiments may be carried out by a personal assistant application or service, which may be implemented as one or more computer applications, services, or routines, such as an app running on a mobile device or the cloud, as further described herein.

[0032] Turning now to FIG. 1, a block diagram is provided showing an example operating environment 100 in which some embodiments of the present disclosure may be employed. It should be understood that this and other arrangements described herein are set forth only as examples. Other arrangements and elements (e.g., machines, interfaces, functions, orders, and groupings of functions, etc.) can be used in addition to or instead of those shown, and some elements may be omitted altogether for the sake of clarity. Further, many of the elements described herein are functional entities that may be implemented as discrete or distributed components or in conjunction with other components, and in any suitable combination and location. Various functions described herein as being performed by one or more entities may be carried out by hardware, firmware, and/or software. For instance, some functions may be carried out by a processor executing instructions stored in memory.

[0033] Among other components not shown, example operating environment 100 includes a number of user devices, such as user devices 102a and 102b through 102n; a number of data sources, such as data sources 104a and 104b through 104n; server 106; and network 110. It should be understood that environment 100 shown in FIG. 1 is an example of one suitable operating environment. Each of the components shown in FIG. 1 may be implemented via any type of computing device, such as computing device 600 described in connection to FIG. 6, for example. These components may communicate with each other via network 110, which may include, without limitation, one or more local area networks (LANs) and/or wide area networks (WANs). In exemplary implementations, network 110 comprises the Internet and/or a cellular network, amongst any of a variety of possible public and/or private networks.

[0034] It should be understood that any number of user devices, servers, and data sources may be employed within operating environment 100 within the scope of the present disclosure. Each may comprise a single device or multiple

devices cooperating in a distributed environment. For instance, server 106 may be provided via multiple devices arranged in a distributed environment that collectively provide the functionality described herein. Additionally, other components not shown may also be included within the distributed environment.

[0035] User devices 102a and 102b through 102n may be client devices on the client-side of operating environment 100, while server 106 may be on the server-side of operating environment 100. Server 106 can comprise server-side software designed to work in conjunction with client-side software on user devices 102a and 102b through 102n so as to implement any combination of the features and functionalities discussed in the present disclosure. This division of operating environment 100 is provided to illustrate one example of a suitable environment, and there is no requirement for each implementation that any combination of server 106 and user devices 102a and 102b through 102n remain as separate entities.

[0036] User devices 102a and 102b through 102n may comprise any type of computing device capable of use by a user. For example, in one embodiment, user devices 102a through 102n may be the type of computing device described in relation to FIG. 6 herein. By way of example and not limitation, a user device may be embodied as a personal computer (PC), a laptop computer, a mobile or mobile device, a smartphone, a tablet computer, a smart watch, a wearable computer, a personal digital assistant (PDA), an MP3 player, a global positioning system (GPS) or device, a video player, a handheld communications device, a gaming device or system, an entertainment system, a vehicle computer system, an embedded system controller, a camera, a remote control, a bar code scanner, a computerized measuring device, an appliance, a consumer electronic device, a workstation, or any combination of these delineated devices, or any other suitable device.

[0037] Data sources 104a and 104b through 104n may comprise data sources and/or data systems, which are configured to make data available to any of the various constituents of operating environment 100, or system 200 described in connection to FIG. 2. (For example, in one embodiment, one or more data sources 104a through 104n provide (or make available for accessing) user data to user-data collection component 210 of FIG. 2.) Data sources 104a and 104b through 104n may be discrete from user devices 102a and 102b through 102n and server 106, or may be incorporated and/or integrated into at least one of those components. In one embodiment, one or more of data sources 104a through 104n comprise one or more sensors, which may be integrated into or associated with one or more of the user device(s) 102a, 102b, or 102n or server 106. Examples of sensed user data made available by data sources 104a through 104n are described further in connection to user-data collection component 210 of FIG. 2.

[0038] Operating environment 100 can be utilized to implement one or more of the components of system 200, described in FIG. 2, and system 300, described in FIG. 3, including components for collecting user data, monitoring activity events, determining activity patterns, consuming activity pattern information to provide an improved user experience, generating personalized content, and/or presenting notifications and related content to users. Referring now to FIG. 2, with FIG. 1, a block diagram is provided showing aspects of an example computing system architecture suit-

able for implementing an embodiment and designated generally as system **200**. System **200** represents only one example of a suitable computing system architecture. Other arrangements and elements can be used in addition to or instead of those shown, and some elements may be omitted altogether for the sake of clarity. Further, as with operating environment **100**, many of the elements described herein are functional entities that may be implemented as discrete or distributed components or in conjunction with other components, and in any suitable combination and location.

[0039] Example system **200** includes network **110**, which is described in connection to FIG. **1**, and which communicatively couples components of system **200** including user-data collection component **210**, presentation component **220**, user activity monitor **280**, activity pattern inference engine **260**, activity pattern consumers **270**, and storage **225**. User activity monitor **280** (including its components **282**, **284**, and **286**), activity pattern inference engine **260** (including its components **262**, **264**, **266**, and **267**), user-data collection component **210**, presentation component **220**, and activity pattern consumers **270** may be embodied as a set of compiled computer instructions or functions, program modules, computer software services, or an arrangement of processes carried out on one or more computer systems, such as computing device **600** described in connection to FIG. **6**, for example.

[0040] In one embodiment, the functions performed by components of system **200** are associated with one or more personal assistant applications, services, or routines. In particular, such applications, services, or routines may operate on one or more user devices (such as user device **102a**), servers (such as server **106**), may be distributed across one or more user devices and servers, or be implemented in the cloud. Moreover, in some embodiments, these components of system **200** may be distributed across a network, including one or more servers (such as server **106**) and client devices (such as user device **102a**), in the cloud, or may reside on a user device, such as user device **102a**. Moreover, these components, functions performed by these components, or services carried out by these components may be implemented at appropriate abstraction layer(s) such as the operating system layer, application layer, hardware layer, etc., of the computing system(s). Alternatively, or in addition, the functionality of these components and/or the embodiments described herein can be performed, at least in part, by one or more hardware logic components. For example, and without limitation, illustrative types of hardware logic components that can be used include Field-programmable Gate Arrays (FPGAs), Application-specific Integrated Circuits (ASICs), Application-specific Standard Products (ASSPs), System-on-a-chip systems (SOCs), Complex Programmable Logic Devices (CPLDs), etc. Additionally, although functionality is described herein with regards to specific components shown in example system **200**, it is contemplated that in some embodiments functionality of these components can be shared or distributed across other components.

[0041] Continuing with FIG. **2**, user-data collection component **210** is generally responsible for accessing or receiving (and in some cases also identifying) user data from one or more data sources, such as data sources **104a** and **104b** through **104n** of FIG. **1**. In some embodiments, user-data collection component **210** may be employed to facilitate the accumulation of user data of a particular user (or in some

cases, a plurality of users including crowdsourced data) for user activity monitor **280**, activity pattern inference engine **260**, or an activity pattern consumer **270**. The data may be received (or accessed), and optionally accumulated, reformatted, and/or combined, by user-data collection component **210** and stored in one or more data stores such as storage **225**, where it may be available to other components of system **200**. For example, the user data may be stored in or associated with a user profile **240**, as described herein. In some embodiments, any personally identifying data (i.e., user data that specifically identifies particular users) is either not uploaded or otherwise provided from the one or more data sources with user data, is not permanently stored, and/or is not made available to user activity monitor **280** and/or activity pattern inference engine **260**.

[0042] User data may be received from a variety of sources where the data may be available in a variety of formats. For example, in some embodiments, user data received via user-data collection component **210** may be determined via one or more sensors, which may be on or associated with one or more user devices (such as user device **102a**), servers (such as server **106**), and/or other computing devices. As used herein, a sensor may include a function, routine, component, or combination thereof for sensing, detecting, or otherwise obtaining information such as user data from a data source **104a**, and may be embodied as hardware, software, or both. By way of example and not limitation, user data may include data that is sensed or determined from one or more sensors (referred to herein as sensor data), such as location information of mobile device (s), properties or characteristics of the user device(s) (such as device state, charging data, date/time, or other information derived from a user device such as a mobile device), user-activity information (for example: app usage; online activity; searches; voice data such as automatic speech recognition; activity logs; communications data including calls, texts, instant messages, and emails; website posts; other user data associated with communication events; etc.) including, in some embodiments, user activity that occurs over more than one user device, user history, session logs, application data, contacts data, calendar and schedule data, notification data, social-network data, news (including popular or trending items on search engines or social networks), online gaming data, ecommerce activity (including data from online accounts such as Microsoft®, Amazon.com®, Google®, eBay®, PayPal®, video-streaming services, gaming services, or Xbox Live®), user-account(s) data (which may include data from user preferences or settings associated with a personal assistant application or service), home-sensor data, appliance data, global positioning system (GPS) data, vehicle signal data, traffic data, weather data (including forecasts), wearable device data, other user device data (which may include device settings, profiles, network-related information (e.g., network name or ID, domain information, workgroup information, connection data, Wi-Fi network data, or configuration data, data regarding the model number, firmware, or equipment, device pairings, such as where a user has a mobile phone paired with a Bluetooth headset, for example, or other network-related information)), gyroscope data, accelerometer data, payment or credit card usage data (which may include information from a user's PayPal account), purchase history data (such as information from a user's Xbox Live, Amazon.com, or eBay account), other sensor data that may be sensed

or otherwise detected by a sensor (or other detector) component(s) including data derived from a sensor component associated with the user (including location, motion, orientation, position, user-access, user-activity, network-access, user-device-charging, or other data that is capable of being provided by one or more sensor components), data derived based on other data (for example, location data that can be derived from Wi-Fi, Cellular network, or IP address data), and nearly any other source of data that may be sensed or determined as described herein.

[0043] In some respects, user data may be provided in user-data streams or signals. A “user signal” can be a feed or stream of user data from a corresponding data source. For example, a user signal could be from a smartphone, a home-sensor device, a GPS device (e.g., for location coordinates), a vehicle-sensor device, a wearable device, a user device, a gyroscope sensor, an accelerometer sensor, a calendar service, an email account, a credit card account, or other data sources. In some embodiments, user-data collection component 210 receives or accesses data continuously, periodically, or as needed.

[0044] User activity monitor 280 is generally responsible for monitoring user data for information that may be used for determining user activity information, which may include identifying and/or tracking features (sometimes referred to herein as “variables”) or other information regarding specific user actions and related contextual information. Embodiments of user activity monitor 280 may determine, from the monitored user data, user activity associated with a particular user. As described previously, the user activity information determined by user activity monitor 280 may include user activity information from multiple user devices associated with the user and/or from cloud-based services associated with the user (such as email, calendars, social-media, or similar information sources), and which may include contextual information associated with the identified user activity. User activity monitor 280 may determine current or near-real-time user activity information and may also determine historical user activity information, in some embodiments, which may be determined based on gathering observations of user activity over time, accessing user logs of past activity (such as browsing history, for example). Further, in some embodiments, user activity monitor 280 may determine user activity (which may include historical activity) from other similar users (i.e., crowdsourcing), as described previously.

[0045] In some embodiments, information determined by user activity monitor 280 may be provided to activity pattern inference engine 260 including information regarding the current context and historical visits (historical observations). Some embodiments may also provide user activity information, such as current user activity, to one or more activity pattern consumers 270. As described previously, user activity features may be determined by monitoring user data received from user-data collection component 210. In some embodiments, the user data and/or information about the user activity determined from the user data is stored in a user profile, such as user profile 240.

[0046] In an embodiment, user activity monitor 280 comprises one or more applications or services that analyze information detected via one or more user devices used by the user and/or cloud-based services associated with the user, to determine activity information and related contextual information. Information about user devices associated

with a user may be determined from the user data made available via user-data collection component 210, and may be provided to user activity monitor 280, activity pattern inference engine 260, or other components of system 200.

[0047] More specifically, in some implementations of user activity monitor 280, a user device may be identified by detecting and analyzing characteristics of the user device, such as device hardware, software such as operating system (OS), network-related characteristics, user accounts accessed via the device, and similar characteristics. For example, information about a user device may be determined using functionality of many operating systems to provide information about the hardware, OS version, network connection information, installed application, or the like.

[0048] Some embodiments of user activity monitor 280, or its subcomponents, may determine a device name or identification (device ID) for each device associated with a user. This information about the identified user devices associated with a user may be stored in a user profile associated with the user, such as in user accounts and devices 244 of user profile 240. In an embodiment, the user devices may be polled, interrogated, or otherwise analyzed to determine information about the devices. This information may be used for determining a label or identification of the device (e.g., a device ID) so that user interaction with the device may be recognized from user data by user activity monitor 280. In some embodiments, users may declare or register a device, such as by logging into an account via the device, installing an application on the device, connecting to an online service that interrogates the device, or otherwise providing information about the device to an application or service. In some embodiments, devices that sign into an account associated with the user, such as a Microsoft® account or Net Passport, email account, social network, or the like, are identified and determined to be associated with the user.

[0049] As shown in example system 200, user activity monitor 280 comprises a user activity detector 282, contextual information extractor 284, and an activity features determiner 286. In some embodiments, user activity monitor 280, one or more of its subcomponents, or other components of system 200, such as activity pattern consumers 270 or activity pattern inference engine 260, may determine interpretive data from received user data. Interpretive data corresponds to data utilized by these components of system 200 or subcomponents of user activity monitor 280 to interpret user data. For example, interpretive data can be used to provide other context to user data, which can support determinations or inferences made by the components or subcomponents. Moreover, it is contemplated that embodiments of user activity monitor 280, its subcomponents, and other components of system 200 may use user data and/or user data in combination with interpretive data for carrying out the objectives of the subcomponents described herein. Additionally, although several examples of how user activity monitor 280 and its subcomponents may identify user activity information are described herein, many variations of user activity identification and user activity monitoring are possible in various embodiments of the disclosure.

[0050] User activity detector 282, in general, is responsible for determining (or identifying) a user action or activity event has occurred. Embodiments of activity detector 282 may be used for determining current user activity or one or more historical user actions. Some embodiments of activity

detector **282** may monitor user data for activity-related features or variables corresponding to user activity such as indications of applications launched or accessed, files accessed, modified, copied, etc., websites navigated to, online content downloaded and rendered or played, or similar user activities.

[0051] Additionally, some embodiments of user activity detector **282** extract from the user data information about user activity, which may include current user activity, historical user activity, and/or related information such as contextual information. (Alternatively or in addition, in some embodiments, contextual information extractor **284** determines and extracts contextual information. Similarly, in some embodiments, activity features determiner **286** extracts information about user activity, such as user activity related features, based on an identification of the activity determined by user activity detector **282**.) Examples of extracted user activity information may include app usage, online activity, searches, calls, usage duration, application data (e.g., emails, messages, posts, user status, notifications, etc.), or nearly any other data related to user interactions with the user device or user activity via a user device. Among other components of system **200**, the extracted user activity information determined by user activity detector **282** may be provided to other subcomponents of user activity monitor **280**, activity pattern inference engine **260**, or one or more activity pattern consumers **270**. Further, the extracted user activity may be stored in a user profile associated with the user, such as in user activity information component **242** of user profile **240**. In some embodiments, user activity detector **282** or user activity monitor **280** (or its other sub components) performs conflation on the detected user activity information. For example, overlapping information may be merged and duplicated or redundant information eliminated.

[0052] In some embodiments, the user activity-related features may be interpreted to determine a user activity has occurred. For example, in some embodiments, user activity detector **282** employs user activity event logic, which may include rules, conditions, associations, classification models, or other criteria, to identify user activity. For example, in one embodiment, user activity event logic may include comparing user activity criteria with the user data in order to determine that an activity event has occurred. The activity event logic can take many different forms depending on the mechanism used to identify an activity event. For example, the user activity event logic could be training data used to train a neural network that is used to evaluate user data to determine when an activity event has occurred. The activity event logic may comprise fuzzy logic, neural network, finite state machine, support vector machine, logistic regression, clustering, or machine-learning techniques, similar statistical classification processes, or combinations of these to identify activity events from user data. For example, activity event logic may specify types of user device interaction(s) information that are associated with an activity event, such as navigating to a website, composing an email, or launching an app. In some embodiments, a series or sequence of user device interactions may be mapped to an activity event, such that the activity event may be detected upon determining that the user data indicates the series or sequence of user interactions has been carried out by the user.

[0053] In some embodiments, activity event logic may specify types of user-device related activity that is consid-

ered user activity, such as activity that happens while a user is logged into the user device, while user interfaces are receiving input (e.g., while a computer mouse, touchpad, screen, voice-recognition interface, or the like are active), or certain types of activity like launching applications, modifying files with applications, opening a browser and navigating to a website, etc. In this way, the activity event logic may be used to distinguish genuine user activity from automated activity of processes running on the user devices, such as automatic updates or malware scanning. Once a user activity is determined, these features or additional related features may be detected and associated with the detected activity for use in determining activity patterns.

[0054] In some embodiments, user activity detector **282** runs on or in association with each user device for a user. User activity detector **282** may include functionality that polls or analyzes aspects of the operating system to determine user activity related features (such as installed or running applications or file accesses and modifications, for example), network communications, and/or other user actions detectable via the user device including sequences of actions.

[0055] In some embodiments, such as the embodiment shown in system **200**, user activity detector **282** includes subcomponents comprising an app activity logging pipeline **283** and a browse activity logging pipeline **285**. These logging pipelines may be embodied as client-side applications or services that run on each user device associated with a user, and in some embodiments may run in conjunction with applications or inside (or as a part of) applications, such as within a browser or as a browser plug-in or extension. App activity logging pipeline **283**, in general, manages logging of a user's application (or app) activity, such as application download, launch, access, use (which may include duration), file access via the application, and in-application user activity (which may include application content). Browse activity logging pipeline **285**, in general, manages logging of a user's browse activity, such as websites visited, social media activity (which may include browse-type activity conducted via specific browsers or apps like the Facebook® app, Twitter® app, Instagram® app, Pinterest® app, etc.), content downloaded, files accessed, and other browse-related user activity. In some embodiments, each browser on a user device is associated with an instance of browse activity logging pipeline **285**, or alternatively a plug-in or service that provided browse information to a single instance of browse activity logging pipeline **285** on the user device. In some embodiments, app and browse activity logging pipelines **283** and **285** may also perform functionality described in connection with contextual information extractor **284**, such as logging time stamps, location stamps, user-device related information, or other contextual information that is associated with the logged app activity or browse activity. In some embodiments, app and browse activity logging pipelines **283** and **285** upload logged user activity information to activity pattern inference engine **260** and/or store the logged activity information in a user profile associated with the user, such as in user activity information component **242** of user profile **240**.

[0056] Contextual information extractor **284**, in general, is responsible for determining contextual information related to the user activity (detected by user activity detector **282** or user activity monitor **280**), such as context features or variables associated with user activity, related information,

and user-related activity, and further responsible for associating the determined contextual information with the detected user activity. In some embodiments, contextual information extractor **284** may associate the determined contextual information with the related user activity and may also log the contextual information with the associated user activity. Alternatively, the association or logging may be carried out by another service. For example, some embodiments of contextual information extractor **284** provide the determined contextual information to activity features determiner **286**, which determines activity features of the user activity and/or related contextual information.

[0057] Some embodiments of contextual information extractor **284** determine contextual information related to a user action or activity event such as entities identified in a user activity or related to the activity (e.g., recipients of a group email sent by the user), user activity associated with the location or venue of the user's device. By way of example and not limitation, this may include context features such as location data; which may be represented as a location stamp associated with the activity; contextual information about the location, such as venue information (e.g., this is the user's office location, home location, school, restaurant, movie theater, etc.), yellow pages identifier (YPID) information, time, day, and/or date, which may be represented as a time stamp associated with the activity; user device characteristics or user device identification information regarding the device on which the user carried out the activity; duration of the user activity, other user activity/activities preceding and/or following the user activity (such as sequences of websites visited, a sequence of online searches conducted, sequences of application and website usage, such as browsing to a bank and then accessing an Excel® spreadsheet file to record financial information, or the like), other information about the activity such as entities associated with the activity (e.g., venues, people, objects, etc.), information detected by sensor(s) on user devices associated with the user that is concurrent or substantially concurrent to the user activity (e.g., motion information or physiological information detected on a fitness tracking user device, listening to music, which may be detected via a microphone sensor if the source of the music is not a user device), or any other information related to the user activity that is detectable that may be used for determining patterns of user activity.

[0058] In embodiments using contextual information related to user devices, a user device may be identified by detecting and analyzing characteristics of the user device, such as device hardware, software such as operating system (OS), network-related characteristics, user accounts accessed via the device, and similar characteristics. For example, as described previously, information about a user device may be determined using functionality of many operating systems to provide information about the hardware, OS version, network connection information, installed application, or the like. In some embodiments, a device name or identification (device ID) may be determined for each device associated with a user. This information about the identified user devices associated with a user may be stored in a user profile associated with the user, such as in user account(s) and device(s) **244** of user profile **240**. In an embodiment, the user devices may be polled, interrogated, or otherwise analyzed to determine contextual information about the devices. This information may be used for deter-

mining a label or identification of the device (e.g., a device ID) so that user activity on one user device may be recognized and distinguished from user activity on another user device. Further, as described previously, in some embodiments, users may declare or register a user device, such as by logging into an account via the device, installing an application on the device, connecting to an online service that interrogates the device, or otherwise providing information about the device to an application or service. In some embodiments devices that sign into an account associated with the user, such as a Microsoft® account or Net Passport, email account, social network, or the like, are identified and determined to be associated with the user.

[0059] In some implementations, contextual information extractor **284** may receive user data from user-data collection component **210**, parse the data, in some instances, and identify and extract context features or variables (which may also be carried out by activity features determiner **286**). Context variables may be stored as a related set of contextual information associated with the user activity, and may be stored in a user profile such as in user activity information component **242**. In some cases, contextual information may be used by one or more activity pattern consumers, such as for personalizing content or a user experience, such as when, where, or how to present content. Contextual information also may be determined from the user data of one or more users, in some embodiments, which may be provided by user-data collection component **210** in lieu of or in addition to user activity information for the particular user.

[0060] Activity features determiner **286** is generally responsible for determining activity-related features (or variables) associated with the user activity that may be used for identifying patterns of user activity. Activity features may be determined from information about a user activity and/or from related contextual information. In some embodiments, activity features determiner **286** receives user-activity or related information from user activity monitor **280** (or its subcomponents), and analyzes the received information to determine a set of one or more features associated with the user activity.

[0061] Examples of activity-related features include, without limitation, location-related features, such as location of the user device(s) during the user activity, venue-related information associated with the location, or other location-related information; time-related features, such as time(s) of day(s), day of week or month, or the duration of the activity, or related duration information such as how long the user used an application associated with the activity; user device-related features, such as device type (e.g., desktop, tablet, mobile phone, fitness tracker, heart rate monitor, etc.), hardware properties or profiles, OS or firmware properties, device IDs or model numbers, network-related information (e.g., mac address, network name, IP address, domain, work group, information about other devices detected on the local network, router information, proxy or VPN information, other network connection information, etc.), position/motion/orientation-related information about the user device, power information such as battery level, time of connecting/disconnecting a charger, and user-access/touch information; usage-related features, such as file(s) accessed, app usage (which may also include application data, in-app usage, concurrently running applications), network usage information, user account(s) accessed or otherwise used (such as device account(s), OS level account(s), or online/cloud-

services related account(s) activity, such as Microsoft® account or Net Passport, online storage account(s), email, calendar, or social networking accounts, etc.); content-related features, such as online activity (e.g., searches, browsed websites, purchases, social networking activity, communications sent or received including social media posts); other features that may be detected concurrent with the user activity or near the time or the user activity; or any other features that may be detected or sensed and used for determining a pattern of the user activity. Features may also include information about user(s) using the device; other information identifying a user, such as a login password, biometric data, which may be provided by a fitness tracker or biometric scanner; and/or characteristics of the user(s) who use the device, which may be useful for distinguishing users on devices that are shared by more than one user. In some embodiments, user activity event logic (described in connection to user activity detector 282) may be utilized to identify specific features from user activity information.

[0062] Continuing with system 200 of FIG. 2, activity pattern inference engine 260 is generally responsible for determining user activity patterns based on the user activity information determined from user activity monitor 280. In some embodiments, activity pattern inference engine 260 may run on a server, as a distributed application across multiple devices, or in the cloud. At a high level, activity pattern inference engine 260 may receive user activity-related information, which may be uploaded from user activity logs from client-side applications or services associated with user activity monitor 280. One or more inference algorithms may be applied to the user activity-related information to determine a set of likely user activity patterns. For example, patterns may be determined based on similar instances of observation of user activity or associated contextual information, which may be referred to as “in-common features” of user activity-related information. The inferred activity pattern information may be provided to an activity pattern consumer 270 and/or used to generate a pattern-based prediction regarding likely future user action (s). In some embodiments, a corresponding confidence is also determined for the patterns (or predictions based on the patterns), as described herein. Further, the activity pattern (or prediction of future action based on a pattern) may comprise a single (future-occurring) user activity likely to occur, a sequence of future user actions, or probabilities for more than one future action, for example, an eighty percent likelihood that the next action will be browsing to website A and a fifteen percent likelihood that the next action will be launching a music player application, and a five percent likelihood that the next action will be browsing to website B.

[0063] As shown in example system 200, activity pattern inference engine 260 comprises semantic information analyzer 262, features similarity identifier 264, and activity pattern determiner 266. Semantic information analyzer 262 is generally responsible for determining semantic information associated with the user activity-related features identified by user activity monitor 280. For example, while a user-activity feature may indicate a specific website visited by the user, semantic analysis may determine the category of website, related websites, themes or topics or other entities associated with the website or user activity. Semantic information analyzer 262 may determine additional user activity-related features semantically related to the user activity, which may be used for identifying user activity patterns.

[0064] In particular, as described previously, a semantic analysis is performed on the user activity information, which may include the contextual information, to characterize aspects of the user action or activity event. For example, in some embodiments, activity features associated with an activity event may be classified or categorized (such as by type, time frame or location, work-related, home-related, themes, related entities, other user(s) (such as communication to or from another user) and/or relation of the other user to the user (e.g., family member, close friend, work acquaintance, boss, or the like), or other categories), or related features may be identified for use in determining a similarity or relational proximity to other user activity events, which may indicate a pattern. In some embodiments, semantic information analyzer 262 may utilize a semantic knowledge representation, such as a relational knowledge graph. Semantic information analyzer 262 may also utilize semantic analysis logic, including rules, conditions, or associations to determine semantic information related to the user activity. For example, a user activity event comprising an email sent to someone who works with the user may be characterized as a work-related activity. Thus, where the user emails some person she works with every Sunday night, but not necessarily the same person, a pattern may be determined (using activity pattern determiner 266) that the user performs work-related activities every Sunday night. Accordingly, it may be appropriate to surface a notification, such as a reminder, relating to the user's work to the user on Sunday night since the user has a pattern of working on Sunday night. (Here, a notification service is one example of an activity pattern consumer 270).

[0065] Semantic information analyzer 262 may also be used to characterize contextual information associated with the user activity event, such as determining that a location associated with the activity corresponds to a hub or venue of interest to the user (such as the user's home, work, gym, or the like) based on frequency of user visits. For example, the user's home hub may be determined (using semantic analysis logic) to be the location where the user spends most of her time between 8 PM and 6 AM. Similarly, the semantic analysis may determine time of day that corresponds to working hours, lunch time, commute time, etc. Similarly, the semantic analysis may categorize the activity as being associated with work or home, based on other characteristics of the activity (e.g., a batch of online searches about chi-squared distribution that occurs during working hours at a location corresponding to the user's office may be determined to be work-related activity, whereas streaming a movie on Friday night at a location corresponding to the user's home may be determined to be home-related activity). In this way, the semantic analysis provided by semantic information analyzer 262 may provide other relevant features of the user activity events that may be used for determining user activity patterns. For example, where the user activity comprises visiting CNN.com over lunch, and the semantic analysis determines that the user visited a news-related website over lunch, a pattern of user activity may be determined (by activity pattern determiner 266) indicating that the user routinely visits news-related websites over lunch, but only occasionally visits CNN.com as one of those news-related websites.

[0066] Features similarity identifier 264 is generally responsible for determining similarity of activity features of two or more user activity events (put another way, activity

features characterizing a first user activity event that are similar to activity features characterizing a second user activity event). The activity features may include features relating to contextual information and features determined by semantic information analyzer 262. Activity events having in-common activity features may be used to identify an activity pattern, which may be determined using activity pattern determiner 266.

[0067] For example, in some embodiments, features similarity identifier 264 may be used in conjunction with one or more pattern-based predictors 267 (a subcomponent of activity pattern determiner 266) to determine a set of user activity events that have in-common features. In some embodiments, this set of user activity events may be used as inputs to a pattern-based predictor, as described below. In some embodiments, features similarity identifier 264 comprises functionality for determining similarity of periodic and behavioral-based activity features. Periodic features comprise, for example, features that may occur periodically; for example, features on a day of the week or month, even/odd days (or weeks), monthly, yearly, every other day, every 3rd day, etc. Behavior features may comprise behaviors such as user activities that tend to occur with certain locations or activities occurring before or after a given user activity event (or sequence of previous activity events), for example.

[0068] In embodiments where activity features have a value, similarity may be determined among different activity features having the same value or approximately the same value, based on the particular feature. (For example, a time stamp of a first activity happening at 9:01 AM on Friday and a time stamp of a second activity happening at 9:07 AM on Friday may be determined to have similar or in-common time stamp features.)

[0069] Activity pattern determiner 266 is generally responsible for determining a user activity pattern based on similarities identified in user activity information. In particular, activity pattern determiner 266 (or activity pattern inference engine 260) may determine a user activity pattern based on repetitions of similar activity features associated with a plurality of observed activity events. Thus, for example, an activity pattern may be determined where activity features corresponding to two or more activity events are similar. In some instances, an activity event may have many corresponding activity features, which may be represented as a feature vector associated with a particular activity event. Accordingly, the analysis carried out by activity pattern determiner 266 may involve comparing the activity features from features vectors of a plurality of activity events.

[0070] In some embodiments, activity patterns may be determined using pattern inferences logic 230. Pattern inferences logic may include rules, associations, conditions, prediction and/or classification models, or pattern inference algorithms. The pattern inferences logic 230 can take many different forms depending on the particular activity pattern or the mechanism used to identify an activity pattern, or identify feature similarity among observed activity events to determine the pattern. For example, some embodiments of pattern inferences logic 230 may employ machine-learning mechanisms to determine feature similarity, or other statistical measures to determine the activity events belonging to a set of “example user actions” that support the determined activity pattern, as further described below. The user activity

information may be received from user activity monitor 280 and information about identified similar features may be received from features similarity identifier 264. In some embodiments, the user pattern(s) determined by activity pattern determiner 266 may be stored as inferred user routines 248 in user profile 240.

[0071] In some embodiments, activity pattern determiner 266 provides a pattern of user activity and an associated confidence score regarding the strength of the user pattern, which may reflect the likelihood that future user activity will follow the pattern. More specifically, in some embodiments, a corresponding confidence weight or confidence score may be determined regarding a determined user activity pattern. The confidence score may be based on the strength of the pattern, which may be determined based on the number of observations (of a particular user activity event) used to determine a pattern, how frequently the user's actions are consistent with the activity pattern, the age or freshness of the activity observations, the number of similar features, types of features, and/or degree of similarity of the features in common with the activity observations that make up the pattern, or similar measurements.

[0072] In some instances, the confidence score may be considered when providing a determined activity pattern to an activity pattern consumer 270. For example, in some embodiments, a minimum confidence score may be needed before using the activity pattern to provide an improved user experience or other service by an activity pattern consumer 270. In one embodiment, a threshold of 0.6 (or just over fifty percent) is utilized such that only activity patterns having a 0.6 (or greater) likelihood of predicting user activity may be provided. Nevertheless, where confidence scores and thresholds are used, determined patterns of user activity with confidence scores less than the threshold still may be monitored and updated based on additional activity observations, since the additional observations may increase the confidence for a particular pattern.

[0073] Some embodiments of activity pattern determiner 266 determine a pattern according to the example approaches described below, where each instance of a user activity event has corresponding historical values of tracked activity features (variables) that form patterns, and where activity pattern determiner 266 may evaluate the distribution of the tracked variables for patterns. In the following example, a tracked variable for a user activity event is a time stamp corresponding to an observed instance of the user activity event. However, it will be appreciated that, conceptually, the following can be applied to different types of historical values for tracked activity features (variables).

[0074] A bag of time stamps (i.e., values of a given tracked variable) can be denoted as $\{t_m\}_{m=1}^M$, and mapped to a two-dimensional histogram of hours and days of the week. The two-dimensional histogram can comprise a summation over the instances of the user-device interaction, such as:

$$h_{ij} = \sum_{m=1}^M I[\text{dayOfWeek}[t_m] = i] \wedge I[\text{hourOfDay}[t_m] = j].$$

This histogram can be used to determine derivative histograms. For example, a day of the week histogram may correspond to: $h_i = \sum_j h_{ij}$. An hour of the day histogram may correspond to: $h_j = \sum_i h_{ij}$. As further examples, one or more histograms may be determined for particular semantic time

resolutions in the form of: $h_{iC} = \sum_{j \in C} h_{ij}$. Any of various semantic time resolutions may be employed, such as week-days and weekends, or morning, afternoon, and night. An example of the latter is where $C \in \{\text{morning, afternoon, night}\}$, morning= $\{9, 10, 11\}$, afternoon= $\{12, 13, 14, 15, 16\}$, and night= $\{21, 22, 23, 24\}$.

[0075] An additional data structure utilized in representing an event can comprise the number of distinct time stamps in every calendar week that has at least one time stamp therein, which may be represented as:

$$w_i^j = \|\{m \mid t_m \text{ is within the } i\text{-th } j \text{ week period}\}\|.$$

As an example, w_2^3 can denote the number of distinct time stamps during the 2nd three-week period of available time stamps. $N^{(j)}$ may be utilized to denote the number of j-week time stamps available in the tracked data; for example, $N^{(3)}$ denotes the number of three-week periods available in the time stamps.

[0076] Activity pattern determiner **266** (or activity pattern inference engine **260**) may generate a confidence score that quantifies a level of certainty that a particular pattern is formed by the historical values in the tracked variable. In the following example, the above principles are applied utilizing Bayesian statistics. In some implementations, a confidence score can be generated for a corresponding tracked variable that is indexed by a temporal interval of varying resolution. For time stamps, examples include Tuesday at 9 AM, a weekday morning, and a Wednesday afternoon. The confidence score may be computed by applying a Dirichlet-multinomial model and computing the posterior predictive distribution of each period histogram. In doing so, a prediction for each bin in a particular histogram may be given by:

$$x_i = \frac{\alpha_0 + h_i}{\sum_i^K (\alpha_0 + h_i)};$$

where K denotes the number of bins, α_0 is a parameter encoding the strength of prior knowledge, and

$$i^* = \arg \max_i x_i.$$

Then, the pattern prediction is the bin of the histogram corresponding to i^* and its confidence is given by x_{i^*} . As an example, consider a histogram in which morning=3, afternoon=4, and evening=3. Using $\alpha_0=10$, the pattern prediction is afternoon, and the confidence score is

$$\frac{10+4}{(10+3)+(10+4)+(10+3)} = \frac{14}{40} \approx 0.35.$$

In accordance with various implementations, more observations result in an increased confidence score, indicating an increased confidence in the prediction. As an example, consider a histogram in which morning=3000, afternoon=4000, and evening=3000. Using a similar calculation, the confidence score is

$$\frac{4010}{10030} \approx 0.4.$$

[0077] Also, in some implementations, a confidence score can be generated for a corresponding tracked variable that is indexed by a period and a number of time stamps. Examples include 1 visit per week, and 3 visits every 2 weeks. Using a Gaussian posterior, a confidence score may be generated for a pattern for every period resolution, denoted as j. This may be accomplished by employing the formula:

$$\hat{\mu}^{(j)} = \lambda \left(\frac{1}{N^{(j)}} \sum_i w_i^{(j)} \right) + (1 - \lambda) \mu_0, \text{ where } \lambda = \frac{\sigma_0^2}{\frac{\sigma^2}{N^{(j)}} + \sigma_0^2}.$$

In the foregoing, σ^2 is the sample variance, and σ_0^2 and μ_0 are parameters to the formula. A confidence score can be computed by taking a fixed interval around the number of time stamps prediction and computing the cumulative density as:

$$\text{conf}_i = P(|x - \hat{\mu}^{(j)}| < a) = \int_{\hat{\mu}^{(j)} - a}^{\hat{\mu}^{(j)} + a} \mathcal{N}(x \mid \hat{\mu}^{(j)}, \hat{\sigma}^{(j)}),$$

$$\text{where } \hat{\sigma}^{(j)} = \frac{1}{\frac{N^{(j)}}{\sigma^2} + \frac{1}{\sigma_0^2}}.$$

[0078] As an example, consider the following observations: $w_1^{(1)}=10$, $w_2^{(1)}=1$, $w_3^{(1)}=10$, $w_4^{(1)}=0$, $w_1^{(2)}=11$, and $w_2^{(2)}=10$. $N^{(1)}=4$ and $N^{(2)}=2$. Using $\mu_0=1$ and $\sigma_0^2=10$, $\mu^{(1)}=4.075$, and $\text{conf}_1=0.25$. Furthermore, $\mu^{(2)}=10.31$ and $\text{conf}_2=0.99$. In the foregoing example, although fewer time stamps are available for two week periods, the reduced variance in the user signals results in an increased confidence that a pattern exists.

[0079] Having determined that a pattern exists, or that the confidence score for a pattern is sufficiently high (e.g., satisfies a threshold value), activity pattern determiner **266** may identify that a plurality of user activities corresponds to a user activity pattern for the user. As a further example, activity pattern determiner **266** may determine that a user activity pattern is likely to be followed by a user where one or more of the confidence scores for one or more tracked variables satisfy a threshold value.

[0080] In some embodiments, patterns of user activity may be determined by monitoring one or more activity features, as described previously. These monitored activity features may be determined from the user data described previously as tracked variables or as described in connection to user-data collection component **210**. In some cases, the variables can represent context similarities and/or semantic similarities among multiple user actions (activity events). In this way, patterns may be identified by detecting variables or features in common over multiple user actions. More specifically, features associated with a first user action may be correlated with features of a second user action to determine a likely pattern. An identified feature pattern may become stronger (i.e., more likely or more predictable) the more often the user activity observations that make up the pattern

are repeated. Similarly, specific features can become more strongly associated with a user activity pattern as they are repeated.

[0081] In some embodiments, such as the example embodiment shown in system **200**, activity pattern determiner **266** includes one or more pattern-based predictors **267**. Pattern-based predictors **267** comprise one or more predictors for predicting a next or future user action taken by the user based on patterns, such as patterns of behavior or similarity features. At a high level, a pattern-based predictor **267** receives user activity information and/or activity features associated with a user activity and determines a prediction of the next action or a future action taken by the user. In an embodiment, a pattern-based predictor **267** includes functionality, as further described below, for performing user activity filtering, determining an activity score, selecting an activity based on the score, and determining a particular pattern-based prediction.

[0082] In one embodiment, a pattern-based predictor **267** uses features similarity identifier **264** to determine features or patterns in common between historical user activities and a recent user activity. For example, similarity of periodic features may be determined, from among the set of historical user actions, from those historical actions having a periodic feature in common with a current or recent user action. Thus, for example, if a recent user action happened on a Monday, on the first day of the month, on an even week, or on a weekday, then determining periodic features similarity would comprise identifying those historical user actions that have features indicating the user action happened on a Monday, those historical user actions having features corresponding to first day of the month (any first day, not just Mondays), or happening on an even week, or a weekday. Likewise, behavior features similarity may be determined to identify sets of historical user actions having a particular behavior feature in common with a current or recent user action. For example, if a recent user action includes visiting a news-related website, then determining behavior features similarity would comprise identifying those historical user actions that have features indicating the user visited a news-related website.

[0083] User activity filtering may use the feature similarity determinations to filter out historical user actions and retain only those historical user actions that have a particular feature (or features) in common with a current or recent user action. Thus, in some embodiments, each pattern-based predictor **267** may be designed (or tuned) for determining a prediction based on a particular feature (or features); for example, there might be a subset of pattern-based predictors **267** used for determining a prediction when the feature indicates a workday, or weekend, or Monday, or particular app access, or application usage duration, etc. Such a pattern-based predictor **267** may need only those historical user actions corresponding to its prediction model. (In some embodiments, such pattern-based predictors may utilize specific prediction algorithms or models, based on their type of pattern prediction (prediction model). These algorithms or models may be stored in with pattern inferences logic **230** in storage **225**.)

[0084] Accordingly, in some embodiments, for each pattern-based predictor **267**, user activity filtering may be carried out to determine a set of historical user actions that are relevant to that particular pattern-based predictor **267**, which may include, for example, periodic-feature based

predictors, behavior-feature based predictors (which may include behavior sequences or sequences of previous user actions), unique or uncommon behavior features (such as when a user performs an activity at an unusual time when compared to similar historical user activities) or other types of feature-based predictors. (In some embodiments, features similarity identifier **264** may determine user activity sequence similarity (e.g., the sequence of the last K user actions prior to the current action (or a particular recent action) by determining a Levenshtein distance between the historical user action sequence and recent user action sequences).

[0085] Additionally, in embodiments where user activity filtering may be performed such that each predictor **267** determines a subset of historical user actions with features that correspond to its prediction criteria, for each predictor **267**, for the subset of historical user actions that pass the filter, user action scoring may be performed on the subset of features. User action scoring generally compares similarities of features in the current or recent user action and the subset of historical user actions (which may be considered as a comparison of contexts) and scores each user action with respect to the similarity of its features. In particular, some embodiments score not only those features used for determining the subset of historical user actions, but all (or a larger number) of features available in the current or recent action and historical actions for comparison. In some embodiments, a Boolean logic process is used (i.e., the features have to be true or have the same pattern, and if this is satisfied, then the difference between the particular features is determined). The differences may include, for example, differences in the time-related features, usage duration, sequence distances, etc. In an embodiment, these differences are determined and put into a sigmoid. Further, in an embodiment, a similarity threshold is used, which may be predetermined, tunable, or adaptive, or may be initially set to a value based on a population of users or may be based on empirical information learned about the particular user, for example, and may be adaptive based on the number of historical observations. The similarity threshold may be used to determine whether a particular historical user action is “similar enough” to the particular current or recent user action so as to be considered for determining a prediction. In some embodiments, a vector representing the similarity differences (or similarity score) may be determined, for example, where multiple features are evaluated for similarity.

[0086] The user actions from the subset of historical user actions that are most similar (or similar enough) to the particular current or recent user action may be selected. In some embodiments, the selection process uses a similarity threshold, such as described above, to determine those historical user actions that satisfy the similarity threshold. (Although the term “selection” is used, it is contemplated that the selection is performed by a computer-related process that does not require a person to perform the selection.) The selected historical user actions comprise a set of “example user actions.”

[0087] A prediction of the user’s next action (or future action) may be inferred based on (or according to) the historical user actions in the set of example user actions. In an embodiment, the predicted next user action (or future action) is the next user action with the highest observations count (i.e., the next user action that is predicted the most

based on the set of example user actions). Those historical user actions in the set of example user actions that are consistent with the prediction comprise a “prediction support set.” In some embodiments, a prediction probability corresponding to the prediction may be determined, for example, based on a ratio of the size of the prediction support set versus the total number of observations (historical user actions in the subset determined by the user action filtering). Moreover, in some embodiments, the prediction may also comprise additional information related to the predicted user action(s), such as activity features which characterize the predicted action(s). By way of example and not limitation, if the predicted action is that the user will email a work colleague, additional activity features may indicate the specific email recipient, the subject matter of the email, the approximate time the user will access his email program to compose the email, or other information related to the predicted action. The related information may also be determined based on the activity features of the prediction support set observations.

[0088] Some embodiments also determine a prediction significance, which may be determined based on a confidence interval (e.g., a Binomial confidence interval) or other appropriate statistical measure. Still further, in some embodiments, the prediction confidence may be based on the prediction probability and the prediction significance (e.g., the product of the prediction probability and prediction significance). Accordingly, some embodiments of activity pattern determiner **266** provide a predicted next (or future) user action or series of next (or future) actions for each of the pattern-based predictors **267**.

[0089] Some embodiments determine a specific prediction from the one or more pattern-based predictors **267**. In an embodiment, an ensemble process is utilized wherein the one or more pattern-based predictors **267** vote, and a selection is determined based on the ensemble-member predictors. Further, in some embodiments, ensemble-member predictors may be weighted based on learned information about the user actions. In one embodiment, once each pattern-based predictor **267** has provided a prediction, the prediction that has the highest corresponding confidence is determined as the next (or future) predicted user action, and may be considered a pattern-based (or history-based) prediction. In some embodiments, the output of activity pattern determiner **266** may be stored as inferred user routines **248** in user profile **240**, and in some embodiments may be provided to an activity pattern consumer **270**.

[0090] Continuing with FIG. 2, example system **200** includes one or more activity pattern consumers **270**, which comprise applications or services that consume activity pattern information to provide improved user experiences. Examples of activity pattern consumers **270** may include, without limitation, content personalization services, user intent inference services, device functionality services, device power management services, and communications management services.

[0091] In particular, a first example activity pattern consumer **270** comprises content personalization services. In one embodiment, a content personalization engine **271** is provided to facilitate providing a personalized user experience. Thus, content personalization engine **271** may be considered one example of an application or service (or set of applications or services) that may consume information about user activity patterns, which may include the predic-

tions of future user actions as determined by implementations of the present disclosure.

[0092] At a high level, example content personalization engine **271** is responsible for generating and providing aspects of personalized user experiences, such as personalized content or tailored delivery of content to a user. The content may be provided to the user as a personalized notification (such as described in connection with presentation component **220**), may be provided to an application or service of the user (such as a calendar or scheduling application), or may be provided as part of an API where it may be consumed by yet another application or service. In one embodiment, the personalized content includes suggesting that the user perform a relevant activity at the right time before the user performs the activity manually. For example, where an activity pattern indicates the user visits his bank’s website near the beginning of the month and enters financial information into an Excel file, the user may be provided with a recommendation asking the user, “Would you like to visit your bank website?” Upon responding affirmatively, a browser instance may be provided that has automatically navigated to the user’s bank’s website. Further, the particular Excel file may be opened automatically when the user acknowledges the suggestion or when the user navigates to the bank’s website. Still further, the user may be provided this content (here, the notification suggestion) at a convenient time, such as when the user is at his home, on an evening, near the beginning of the month, etc.

[0093] In one embodiment, the personalized content may include a notification which comprises a reminder, a recommendation, suggestion, request, communication-related data (e.g., an email, instant message, or call notification), information relevant to the user in some manner, or other content that is provided to the user in a way that is personalized. For example, content may be provided at a time when the user would most likely desire to receive it, such as a work-related notification provided at a time according to a user pattern indicating that the user is about to begin a work-related activity, and not providing the content at a time when it is likely to be dismissed, ignored, or bothersome.

[0094] In another example of providing personalized content, where an activity pattern indicates that a user typically checks online for coupons when shopping at a particular store, then upon user data indicating that a user has entered the store or is likely going to the store (which may be determined based on a predicted future activity), the user may be automatically provided with the online coupons. Further, the particular coupons provided may be for items that the user typically buys (or for similar items from competitors) based on user activity pattern information, which may include the user’s purchasing habits.

[0095] In some embodiments, content personalization engine **271** tailors content for a user to provide a personalized user experience. For example, content personalization engine **271** may generate a personalized notification to be presented to a user, which may be provided to presentation component **220**. Alternatively, in other embodiments, content personalization engine **271** generates notification content and makes it available to presentation component **220**, which determines when and how (i.e., what format) to present the notification based on user data or user activity pattern information. For example, if a user activity pattern indicates the user is likely to be driving to work at a time when it is relevant to present a particular notification, it may

be appropriate to provide that notification in an audio format, thus personalizing it to the context of the user. In some embodiments, other services or applications operating in conjunction with presentation component 220 determine or facilitate determining when and how to present personalized content. The personalized content may be stored in a user profile 240, such as in a personalized content component 249.

[0096] Some embodiments of content personalization engine 271 evaluate user content to determine how to provide the content in an appropriate manner. For example, content that is determined to be work related may be withheld from presentation to a user until a user activity pattern indicates the user is likely to be conducting work-related activity.

[0097] A second example activity pattern consumer 270 comprises a user intent inference service. A user's intent may be inferred based on user activity pattern information. In particular, whenever a user interacts with a computing device, such as when the user engages in a user activity, it may be assumed that the user has intent. In one embodiment, a user activity pattern may be analyzed along with sensor data collected by a user device regarding a particular user interaction, such as a user request or command. The user's intent may be inferred based on determining a likely intent that is consistent with a user activity pattern. For example, during a user's lunch hour, the user may perform an online search with the intent of locating information about the user's favorite band, R.E.M., and not information about rapid eye movement (REM) sleep. A user activity pattern indicating that the user typically browses music-entertainment related websites over his lunch hour and/or other activity pattern information about the user indicating that the user recently listened to music by R.E.M. may be used to resolve the user's intent for the query, namely search results related to the music group R.E.M. and not results related to rapid eye movement. As another example, suppose while driving home the user issues a spoken request to a user device to "call Pat." However, the user has more than one "Pat" on his contact list, for example, the user's brother is named Pat and several of the user's friends are also named Pat. Based on an inferred pattern of user activity indicating that the user typically calls his family while driving home, it may be determined that the user's intent is to call his brother Pat because this intent is consistent with an activity pattern.

[0098] A third example activity pattern consumer 270 comprises device power management services. For example, in one embodiment of the disclosure, a device power management service is provided that regulates power consumption of a user device based on user activity pattern information. In particular, based on a pattern of when a user typically charges his mobile device, where it is determined that the time until the typical charging time is great enough such that, based on the mean consumption of battery energy, it is projected that the device will run out of battery power before the time that user typically charges the device, then the power management service may implement power-saving measures to reduce battery consumption. For example, the power management service may automatically switch the device to a battery-saving mode. Similarly, if it is determined that the user is located at or will be located at a location where the user has previously charged his user device, which may be determined based on location-con-

textual information associated with the user activity related to charging the user device, then the user may be provided a notification recommending that the user charge his user device.

[0099] A fourth example activity pattern consumer 270 comprises services for modifying functionality of user device(s). In particular, user device operations, user preferences, settings, or configurations associated with a user device may be modified or configured according to information derived from user activity patterns. For example, suppose a user has a pattern of turning down the volume on his smartphone every time he arrives at work in the morning. One embodiment of a service for modifying device functionality may automatically attenuate the volume of the user device based on user data (e.g., sensor data) indicating that the user has arrived at work (or indicating that it is a workday (e.g., a weekday) morning, etc., according to a determined activity pattern). In an embodiment, the user may opt to use the service; for example, the user may be prompted with a question by a personal assistant service, such as "I notice that you turn down the volume on this device whenever you come to this location. Would you like me to automatically handle this for you each time you visit this location?"

[0100] In some embodiments, based on patterns detected from a first device (or a first set of devices) associated with a user, the functionality of other user devices associated with the user may be modified. For example, for the volume limiting service described above, the volume on other devices, such as a user's tablet computer, smartwatch, or second smartphone, may also be adjusted. Further, in some embodiments, the patterns may be determined from crowd-sourced user activity history. Specifically, as described previously, user activity information from a set of users similar to the given user may be used to determine activity patterns consumed by an activity pattern consumer 270 (or content personalization engine 271). For example, in the volume limiting example described above, a pattern of user activity indicating turning down (or up) the volume on user devices upon arriving at (or leaving) a location, such as arriving at work, in the morning, or arriving at a movie theatre, may be determined for a group of users. In this example, the users are similar because they arrive at the same location for work or for a movie. The activity pattern may be imputed to the particular user, who also arrives at this location for work (or arrives at a movie theatre), and from this imputed activity pattern, an inference or prediction may be determined regarding the likelihood that the particular user will turn down (or up) the volume on his user device upon arriving at the location. Again, based on a prediction or inference determined from the imputed pattern (and in some embodiments based further on a confidence score associated with the prediction or inference), the volume may be automatically attenuated (or amplified), or the user may be prompted about whether the user desires a service to automatically handle adjusting the volume upon arriving (or leaving) the location.

[0101] A fifth example activity pattern consumer 270 comprises communications management services. In particular, network bandwidth and other communication-related resources may be more efficiently utilized according to information derived from user activity patterns. For example, suppose a user has a pattern of watching a streaming or downloaded high-definition movie (e.g., a 4k movie)

on Friday nights. One embodiment of a communications management service may cache or buffer the movie in advance of watching it, so that it is available to watch without interruption. (This may be particularly useful for a user who shares network bandwidth with other users who are watching streaming movies at the same time, thereby reducing available network bandwidth.) In some embodiments, the particular movie (or several movies likely to be watched by the user) may be determined from a user's watch list, or may be inferred based on information derived from monitored user activity, such as user tastes, internet searches, movies watched by the user's social media friends, etc. As described previously, some embodiments may use contextual information, such as a user's calendar information, to determine the likelihood that the user will follow the pattern. Thus, where a user's calendar indicates that an event is scheduled for next Friday night, then it may be determined that, because the user is not likely to watch a movie, the movie does not need to be downloaded prior to next Friday night. Further, in some embodiments, where a user has a data cap or where network speed is capped or otherwise limited, the movie may be downloaded at an earlier time so that it is already available to watch on Friday night. Still further, in some embodiments, the user may be prompted in advance of downloading a particular movie about whether the user may be interested in watching the movie.

[0102] Other examples of activity pattern consumers 270 may include, without limitation: (a) a recommendation service that suggests new content to a user based on user patterns and contextual information. For example, a user activity pattern indicates that a user listens to music every Friday night. Contextual information indicates that the user prefers certain bands or styles of music. Accordingly, on a given Friday night, a recommendation is provided to the user to listen to a new artist having a style similar to the user's taste. (b) A user has an activity pattern of going to music concerts to see bands that sing songs that the user frequently enjoys. A personal assistant application service monitors local concerts and determines that a band that sings songs the user listens to is coming to town. The personal assistant application automatically purchases a ticket for the user when the tickets first become available. Alternatively, the personal assistant service checks the user's calendar to determine that the user is available on the date of the concert, and then prompts the user, notifying the user about the concert, and in some embodiments, asking if the user wants the personal assistant service to purchase a ticket. (c) A user has an activity pattern of watching an online movie on Friday nights. A personal assistant service determines that a user reads certain genres of books, based on information about book purchases and/or e-reader activity by the user. Based on the user's taste in books, a movie may be recommended to the user that the user likely will enjoy. The recommended movie may be automatically downloaded in a manner to preserve bandwidth in advance of Friday night.

[0103] Example system 200 also includes a presentation component 220 that is generally responsible for presenting content and related information to a user, such as the personalized content from content personalization engine 271 or content from other activity pattern consumers 270. Presentation component 220 may comprise one or more applications or services on a user device, across multiple user devices, or in the cloud. For example, in one embodiment, presentation component 220 manages the presentation

of content to a user across multiple user devices associated with that user. Based on content logic, device features, associated logical hubs, inferred logical location of the user, and/or other user data, presentation component 220 may determine on which user device(s) content is presented, as well as the context of the presentation, such as how (or in what format and how much content, which can be dependent on the user device or context) it is presented, when it is presented, etc. In particular, in some embodiments, presentation component 220 applies content logic to device features, associated logical hubs, inferred logical locations, or sensed user data to determine aspects of content presentation.

[0104] In some embodiments, presentation component 220 generates user interface features associated with the personalized content. Such features can include interface elements (such as graphics buttons, sliders, menus, audio prompts, alerts, alarms, vibrations, pop-up windows, notification-bar or status-bar items, in-app notifications, or other similar features for interfacing with a user), queries, and prompts.

[0105] As described previously, in some embodiments, a personal assistant service or application operating in conjunction with presentation component 220 determines when and how (e.g., presenting when the user is determined to be at a specific logical location) to present the content. In such embodiments, the content, including content logic, may be understood as a recommendation to the presentation component 220 (and/or personal assistant service or application) for when and how to present the notification, which may be overridden by the personal assistant app or presentation component 220.

[0106] Example system 200 also includes storage 225. Storage 225 generally stores information including data, computer instructions (e.g., software program instructions, routines, or services), logic, profiles, and/or models used in embodiments described herein. In an embodiment, storage 225 comprises a data store (or computer data memory). Further, although depicted as a single data store component, storage 225 may be embodied as one or more data stores or may be in the cloud.

[0107] As shown in example system 200, storage 225 includes activity pattern inferences logic 230, as described previously, and user profiles 240. One example embodiment of a user profile 240 is illustratively provided in FIG. 2. Example user profile 240 includes information associated with a particular user such as user activity information 242, information about user accounts and devices 244, user preferences 246, inferred user routines 248, and personalized content 249. The information stored in user profile 240 may be available to the activity pattern inference engine 260 or other components of example system 200.

[0108] As described previously, user activity information 242 generally includes user information about user actions or activity events, related contextual information, activity features, or other information determined via user activity monitor 280, and may include historical or current user activity information. User accounts and devices 244 generally includes information about user devices accessed, used, or otherwise associated with a user, and/or information related to user accounts associated with the user, for example, online or cloud-based accounts (e.g., email, social media) such as a Microsoft® Net Passport, other accounts such as entertainment or gaming-related accounts (e.g.,

Xbox Live, Netflix, online game subscription accounts, etc.), user data relating to accounts such as user emails, texts, instant messages, calls, other communications, and other content; social network accounts and data, such as news feeds; online activity; and calendars, appointments, application data, other user accounts, or the like. Some embodiments of user accounts and devices 244 may store information across one or more databases, knowledge graphs, or data structures. As described previously, the information stored in user accounts and devices 244 may be determined from user-data collection component 210 or user activity monitor 280 (including one of its subcomponents).

[0109] User preferences 246 generally include user settings or preferences associated with user activity monitoring. By way of example and not limitation, such settings may include user preferences about specific activities (and related information) that the user desires be explicitly monitored or not monitored or categories of activities to be monitored or not monitored, crowdsourcing preferences, such as whether to use crowdsourced information, or whether the user's activity pattern information may be shared as crowdsourcing data; preferences about which activity pattern consumers may consume the user's activity pattern information; and thresholds, and/or notification preferences, as described herein. As described previously, inferred user routines 248 may include one or more user pattern(s) determined by activity pattern determiner 266, and may also include confidence scores associated with the patterns and/or information related to the activity patterns, such as contextual information or semantic information. Personalized content 249 includes personalized content determined from content personalization engine 271, such as pending or scheduled notifications.

[0110] With reference now to FIG. 3, aspects of an example system for determining user activity patterns is provided and referenced generally as system 300. Example system 300 may determine user activity patterns based on browse activity and application activity from across multiple devices. As shown, example system 300 comprises one or more client devices 302, an inferences system 350, and one or more activity pattern consumers 370. In various embodiments, a client device 302 may be embodied as a user device, such as user device 102a described in FIG. 1. Client device 302 includes an app activity logging pipeline 383 and a browse activity logging pipeline 385. These logging pipelines may be embodied as client-side applications or services that run on each user device associated with a user, and in some embodiments may run in conjunction with applications or inside (or as a part of) applications, such as within a browser or as a browser plug-in or extension. App activity logging pipeline 383, in general, manages logging of a user's application (or app) activity, such as application download, launch, access, use (which may include duration), file access via the application, and in-application user activity (which may include application content). Browse activity logging pipeline 385, in general, manages logging of a user's browse activity, such as websites visited, social media activity (which may include browse-type activity conducted via specific browsers or apps like the Facebook® app, Twitter® app, Instagram® app, Pinterest® app, etc.), content downloaded, files accessed, and other browse-related user activity. In some embodiments, app activity logging pipeline 383 and browse activity logging pipeline 385 further include functionality of the embodiments described

in connection with app activity logging pipeline 283 and browse activity logging pipeline 285 of system 200 in FIG. 2.

[0111] As shown in system 300, inferences system 350 comprises activity pattern inference engine 360, which receives app activity and browse activity from one or more client devices 302 and uses that information to infer activity patterns. Activity pattern inference engine 360 may be embodied as an embodiment of activity pattern inference engine 260, described in connection to FIG. 2. Further, some embodiments of inferences system 350 may also include other inference engines (not shown). For example, other inference engines might include components for determining venue visiting inferences (i.e., inferring future venue visits or patterns of venue visits); location inferences (i.e., inferring location for the user when explicit location information is not available); or shadow calendar, which may include inferences about a user's availability based on explicit and inferred scheduled events.

[0112] Inferences system 350 may provide inferred user activity pattern information, including predicted future actions, to an activity pattern consumer 370. In some instances, an activity pattern consumer 370 may be embodied as an activity pattern consumer 270, described in connection to FIG. 2. A further example activity pattern consumers 370 of system 300 includes browser A 372 (or a service associated with browser A) that navigates to a website according to a predicted user activity pattern; for example, automatically navigating to the user's bank website near the beginning of the month, where the user has a pattern of visiting the website at the beginning of the month. Yet another example activity pattern consumers 370 includes app B 373, which may launch and/or load a file, or perform some functionality (e.g., play a song, compose an email, etc.) according to a predicted user activity pattern.

[0113] Turning to FIG. 4, a flow diagram is provided illustrating one example method 400 for inferring a user activity pattern. Each block or step of method 400 and other methods described herein comprises a computing process that may be performed using any combination of hardware, firmware, and/or software. For instance, various functions may be carried out by a processor executing instructions stored in memory. The methods may also be embodied as computer-usable instructions stored on computer storage media. The methods may be provided by a stand-alone application, a service or hosted service (stand-alone or in combination with another hosted service), or a plug-in to another product, to name a few. Accordingly, method 400 may be performed by one or more computing devices, such as a smartphone or other user device, a server, or by a distributed computing platform, such as in the cloud. The activity pattern may be inferred through the analysis of signal data (or user data) gathered from one or more user devices associated with the user.

[0114] At step 410, the method comprises identifying a set of user devices associated with a user. Embodiments of step 410 may determine a set of user devices based on monitoring user data for user-device related information. In some embodiments, the set of user devices identified in step 410 comprises one user device or a plurality of user devices, such as user devices 102a through 102n, described in connection to FIG. 1.

[0115] In one embodiment, information about user devices associated with a user may be determined from the user data

made available via user-data collection component **210**, such as described in connection to user activity monitor **280** of FIG. 2. For example, as described previously, information about a user device may be sensed or otherwise detected from user data, such as by one or more sensors associated with a user device, or may be determined by detecting and analyzing user-device related information in user data to determine characteristics of the user device, such as device hardware, software such as operating system (OS), network-related characteristics, user accounts accessed via the device, and similar characteristics. In one embodiment, the detected user devices (such as user devices **102a** through **102n**) may be polled, interrogated, or otherwise analyzed to determine information about the devices. In some implementations, this information may be used for determining a label or identification of the device (e.g., a device ID) so that the user interaction with one device may be recognized from user interaction on another device. In some embodiments of step **410**, a user device may be identified based on user-provided information, such as the case where a user declares or registers the device, for example, by logging into an account via the user device, installing an application on the device, connecting to an online service that interrogates the device, or otherwise providing information about the device to an application or service.

[0116] At step **420**, the method comprises monitoring the set of user devices to detect a user activity event. Embodiments of step **420** monitor user data associated with the set of user devices identified in step **410**, to identify or detect a user activity event (sometimes referred to herein as a user action). In some instances, an activity event may comprise a series or sequence of user interactions with one or more user devices.

[0117] Some embodiments of step **420** may comprise monitoring sensor data using one or more sensors associated with the set of user devices. Some embodiments of step **420** may use activity event logic to detect the user activity event, as described in connection to user activity detector **282**. Further, some implementations of step **420** may be carried out using a user activity detector component, such as described in system **200** of FIG. 2. Additional details of embodiments of step **420** are provided in connection with user activity detector **282** in FIG. 2.

[0118] At step **430**, the method comprises determining a set of activity features associated with the activity event. Upon detecting or otherwise identifying a user activity event in step **420**, embodiments of step **430** determine a set of activity features associated with the detected activity event. Some embodiments of step **430** determine the set of activity features based at least in part on sensor data (including data that may be derived from sensor data such as interpretive data) provided by one or more sensors associated with the set of user devices. In some embodiments, the sensor data may be provided via a user-data collection component, as described in FIG. 2. In particular, user data related to the detected activity event and which may be determined at least in part from sensor data, and may include interpreted data, contextual data, and/or semantic information related to the detected activity event, is received and analyzed to determine a set of one or more features associated with the user activity.

[0119] As described previously in connection with user activity monitor **280** and activity features determiner **286** of FIG. 2, activity-related features (or variables) associated

with the user activity that may be used for identifying patterns of user activity. The activity features may be determined from information about a user activity, which in some embodiments may include related contextual information about the detected activity. Additional details of determining contextual information related to a detected activity event are described below, and in connection to contextual information extractor **284** of FIG. 2. Further, in some embodiments, the contextual information may include semantic information determined from a semantic analysis performed on the detected activity event and/or one or more activity features associated with the activity event. For example, while a user-activity feature may indicate a specific website visited by the user, semantic analysis may determine the category of website, related websites, themes or topics, or other entities associated with the website or user activity. From the semantic analysis, additional user activity-related features semantically related to the user activity may be determined and used for identifying user activity patterns. Additional details of determining activity features that include semantic information related to the detected activity event are described in connection to semantic information analyzer **262** of FIG. 2.

[0120] Examples of activity-related features may include, without limitation, location-related features, such as location of the user device(s) during the user activity, prior to and/or after the user activity, venue-related information associated with the location, or other location-related information; time-related features, such as time(s) of day(s), day of week or month the user activity, or the duration of the activity, or related duration information such as how long the user used an application associated with the activity; content-related features, such as online activity by the user (e.g., searches, browsed websites, type or category of websites, purchases, social networking activity, communications sent or received including social media posts, which may comprise online activity occurring prior to or after the detected activity event); other features that may be detected concurrent with the user activity or near the time of the user activity; or any other features that may be detected or sensed and used for determining a pattern of the user activity including semantic features that may characterize aspects of the activity, the nature or type of the activity, and/or user interactions making up the activity. Features may also include information about user(s) using the device; and other contextual features such as user device-related features, usage-related features; or other information about the user. By way of example and not limitation, user device-related features may include information about device type (e.g., desktop, tablet, mobile phone, fitness tracker, heart rate monitor, etc.), hardware properties or profiles, OS or firmware properties, device IDs or model numbers, network-related information (e.g., mac address, network name, IP address, domain, work group, information about other devices detected on the local network, router information, proxy or VPN information, other network connection information, etc.), position/motion/orientation-related information about the user device, power information such as battery level, time of connecting/disconnecting a charger, user-access/touch information; usage-related features may include information about file(s) accessed, app usage (which may also include application data, in-app usage, concurrently running applications), network usage information, user account(s) accessed or otherwise used (such as device account(s), OS level account(s), or

online/cloud-services related account(s) activity, such as Microsoft® account or Net Passport, online storage account (s), email, calendar, or social networking accounts, etc.); other information identifying a user may include login password, biometric data, which may be provided by a fitness tracker or biometric scanner; and/or characteristics of the user(s) who use the device, which may be useful for distinguishing users on devices that are shared by more than one user.

[0121] In some embodiments of step 430, user activity event logic (described in connection to user activity detector 282) may be utilized to identify specific features associated with the detected user activity event. Some implementations of step 430 may be carried out using an activity features determiner component, such as described in system 200 of FIG. 2. Additional details of embodiments of step 430 are provided in connection with activity features determiner component 286, in FIG. 2.

[0122] At step 440, a record of the detected activity event and the activity features associated with the event are stored in an activity event data store. The activity event data store may comprise a plurality of records about activity events, wherein each record may include information about a particular activity event including one or more activity features associated with the activity event. In some embodiments, the plurality of records in the activity event data store comprise records of other activity events determined according to steps 410 through 430 of method 400. In some instances, some of the other records may include information about activity events (including associated activity features) derived from other users determined to be similar to the particular user, as described previously. In an embodiment, the activity event data store comprises a user profile, and may be stored in a user activity information component, such as user activity information component 242 of user profile 240 described in connection to FIG. 2.

[0123] At step 450, an activity pattern inference engine is used to identify an activity pattern. In embodiments of step 450, the activity pattern may be detected based on an analysis of at least a portion of the plurality of activity event records to identify a set of activity events having similar activity features. Thus, embodiments of step 450 may comprise determining a set of user activity events that have in-common features, which may be determined using a features similarity identifier component, such as features similarity identifier 264, described in system 200 of FIG. 2.

[0124] In embodiments, the set of activity events having similar activity features may comprise two or more activity events that form the basis of the pattern, as described previously in connection with activity pattern determiner 266. In some embodiments, the similar activity features may be in common with two or more of the activity events, may be substantially similar, or may be similar enough according to a similarity threshold, as described previously. Further, some embodiments may determine a corresponding confidence score associated with the activity pattern. In some instances, the confidence score may indicate a likelihood that the user will behave according to the pattern, and may be used for determining whether to provide activity pattern information to an activity pattern consumer service, such as described in connection to FIG. 2. In some embodiments, one or more pattern based predictors may be utilized, such as described in connection to FIG. 2.

[0125] Step 450 may be carried out by an activity pattern inference engine, such as activity pattern inference engine 260 described in connection to system 200 of FIG. 2, or one of its subcomponents. Additional details of embodiments of step 450 are provided in connection with activity pattern inference engine 260, in FIG. 2.

[0126] At step 460, the activity pattern is stored in an activity pattern data store. In an embodiment, the activity pattern data store comprises a user profile, and may be stored in an inferred user routines component, such as inferred user routines 248 of user profile 240 described in connection to FIG. 2. The activity pattern may be accessed by one or more computing applications that provide enhanced or improved computing experiences for the user, such as providing personalization services (e.g., personalized content or tailoring the delivery of content to the user, as described herein), improved speech recognition, more efficient consumption of bandwidth or power, or improved semantic understanding of the user, which may be used for performing disambiguation or other aspects of understanding input from the user.

[0127] In some embodiments, once an activity pattern is determined, it can be used to determine that a probable future activity event will occur at a future time that is a threshold time from a present time (e.g., at or within three hours, twelve hours, one day, five days, two weeks, etc., from the current time) by analyzing the exercise pattern and/or current or recent information about user activity. The probable future activity event can be associated with a location, time, condition, and/or situation (e.g., the future activity likely occurs following a certain type of event, such as after the user performs another activity), and other contextual data that can be used to facilitate an improved user experience. Further, as described previously, in some embodiments, the user activity pattern, or an inferred user intent or predictions of future activity determined therefrom, may be made available to one or more applications and services that consume this information and provide an improved user experience, such as activity pattern consumers 270, 271, 370, 372, or 373, described in connection to FIGS. 2 and 3. For example, in an embodiment, a modified user experience is provided, which may include presenting (or withholding or delaying) content to the user in a manner that is personalized, such as described in connection to content personalization engine 271 of system 200 (FIG. 2).

[0128] In some embodiments, a cloud system (such as the cloud system described above) and/or a cloud service may be utilized to perform method 400 so as to provide an improved or enhanced user experience (such as personalized content) to multiple services, which may be running on many different user devices. As such, system 200 can save significant processing, bandwidth, storage, and computing resources by centralizing certain functionality. For example, user-data collection component 210 (FIG. 2) can accumulate user data and interpretive data for multiple users or user devices, such that each user device does not require separate and redundant data collection and storage. Additionally, the processing and storage of user profile data can be made more secure by being disassociated from the user device, which is closely tied to the user.

[0129] With reference now to FIG. 5, a flow diagram is provided illustrating one example method 500 for determining a probable future user action based on a pattern of user activity. At step 510, an inferred activity pattern for a user is

accessed. The inferred exercise pattern may be determined as described previously with reference to FIG. 2, 3, or 4. In an embodiment, the activity pattern may be accessed from an activity pattern data store, such as inferred user routines 248 of user profile 240, described in connection to FIG. 2. As mentioned, a user may be associated with multiple activity patterns. A particular pattern can be associated with a confidence level or score that indicates how likely the pattern is to be followed.

[0130] At step 520, a probable future activity event is determined based on the activity pattern. Embodiments of step 530 predict a future activity event, which may include one or a series (or sequence) of future user interactions likely to occur or likely to be desired by the user to transpire. In some embodiments, step 520 may also determine a confidence score associated with the determined probable future activity indicating the likelihood that the activity event will occur or be desired by the user to occur. In some embodiments, the probable future activity may comprise a context, which may be defined by a location, time, user behavior, or other condition, and may be determined based at least in part on sensor data.

[0131] For example, in some instances, the probable future activity event is determined looking at a periodic or behavior context, which may be defined according to periodic features or behavioral features, as described previously. Thus, a periodic context defines when an activity event associated with an activity pattern is likely to occur. For example, if the pattern indicates a user performs a particular activity every Monday, Wednesday, and Friday, then a future activity event could be determined for any Monday, Wednesday, or Friday. Similarly, a behavioral context defines contextual features present when an activity pattern is likely to occur, such as where the pattern indicates a user performs a particular activity when visiting a certain venue, following another activity, or upon another condition manifesting (e.g., a user checks for online coupons every time he visits a particular store). The behavioral context may also be defined in the negative to identify contextual features that define exceptions to the periodic context. For example, the user may not perform certain activities on holidays or when the user is not in a hub or frequently visited venue, such as when the user is traveling or on vacation. In some embodiments wherein a confidence score is determined for the probable future activity, the confidence score may be further determined based on the context.

[0132] In some embodiments, the context may be determined based in part on a current or recently occurring activity event. In particular, in some embodiments of step 520, the probable future activity may also be determined based on a current or recently occurring user activity event. Some embodiments of step 520 may monitor one or more user devices associated with a user to detect a current or recently occurring user activity event. The detected activity event may be used for determining a particular user activity pattern relevant to a context (e.g., the user's current situation) and/or a next or future activity event in the pattern. Accordingly, the recency of the activity event detected in such embodiments of step 520 may be based on a length of time such that a context (e.g., the user's current situation) is still relevant to the activity pattern. Thus, the recency may vary based on the particular activity pattern.

[0133] In some embodiments, the current or recently occurring activity event is determined based on sensor data

from one or more user devices. For example, some embodiments of step 520 may comprise monitoring sensor data using one or more sensors associated with the set of user devices. Some embodiments of step 520 may be performed as described in step 420 of method 400, or may be carried out using user activity detector 282, described in system 200 of FIG. 2.

[0134] At step 530, based on the determined probable future activity, an enhanced user experience may be provided to the user. Embodiments of step 530 provide an enhanced user experience, which may be in the form of a tailored, modified, or personalized service for the user that is determined based at least in part on the probable future activity determined in step 520, and may further be determined based on an associated confidence score. Embodiments of step 530 may comprise providing a service to the user by one or more activity pattern consumers 270, described in connection to FIG. 2. For example, in an embodiment, the enhanced user experience comprises one of a recommendation, notification, request, or suggestion related to the probable future activity. In one embodiment, it comprises automatically carrying out actions related to the probable future activity at a time or location consistent with the inferred activity pattern. For example, where the user typically browses her bank account near the beginning of the month and copies financial information from the bank website into a spreadsheet, an embodiment of step 530 may automatically perform the operations of launching a browser, navigating to the bank website, and loading the spreadsheet file in a spreadsheet application. In some embodiments, the automated operations may be performed at an appropriate time, based on the context, such as while the user is at home using the computing device (user device) that she typically uses to access the bank website, or not while the user is driving and thus unable to benefit from the personalized user experience. Additional examples of enhanced user experiences are described herein and include, by way of example and not limitation, personalization services (e.g., providing personalized content or tailoring the delivery of content to the user, as described herein), improved speech recognition, more efficient consumption of bandwidth or power, or improved semantic understanding of the user, which may be used for performing disambiguation or other aspects of understanding input from the user.

[0135] Having described various implementations, an exemplary computing environment suitable for implementing embodiments of the disclosure is now described. With reference to FIG. 6, an exemplary computing device is provided and referred to generally as computing device 600. The computing device 600 is but one example of a suitable computing environment and is not intended to suggest any limitation as to the scope of use or functionality of embodiments of the disclosure. Neither should the computing device 600 be interpreted as having any dependency or requirement relating to any one or combination of components illustrated.

[0136] Embodiments of the disclosure may be described in the general context of computer code or machine-useable instructions, including computer-useable or computer-executable instructions, such as program modules, being executed by a computer or other machine, such as a personal data assistant, a smartphone, a tablet PC, or other handheld device. Generally, program modules, including routines, programs, objects, components, data structures, and the like,

refer to code that performs particular tasks or implements particular abstract data types. Embodiments of the disclosure may be practiced in a variety of system configurations, including handheld devices, consumer electronics, general-purpose computers, more specialty computing devices, etc. Embodiments of the disclosure may also be practiced in distributed computing environments where tasks are performed by remote-processing devices that are linked through a communications network. In a distributed computing environment, program modules may be located in both local and remote computer storage media including memory storage devices.

[0137] With reference to FIG. 6, computing device 600 includes a bus 610 that directly or indirectly couples the following devices: memory 612, one or more processors 614, one or more presentation components 616, one or more input/output (I/O) ports 618, one or more I/O (I/O) components 620, and an illustrative power supply 622. Bus 610 represents what may be one or more busses (such as an address bus, data bus, or combination thereof). Although the various blocks of FIG. 6 are shown with lines for the sake of clarity, in reality, these blocks represent logical, not necessarily actual, components. For example, one may consider a presentation component such as a display device to be an I/O component. Also, processors have memory. The inventors hereof recognize that such is the nature of the art and reiterate that the diagram of FIG. 6 is merely illustrative of an exemplary computing device that can be used in connection with one or more embodiments of the present disclosure. Distinction is not made between such categories as “workstation,” “server,” “laptop,” “handheld device,” etc., as all are contemplated within the scope of FIG. 6 and with reference to “computing device.”

[0138] Computing device 600 typically includes a variety of computer-readable media. Computer-readable media can be any available media that can be accessed by computing device 600 and includes both volatile and nonvolatile media, removable and non-removable media. By way of example, and not limitation, computer-readable media may comprise computer storage media and communication media. Computer storage media includes both volatile and nonvolatile, removable and non-removable media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, or other data. Computer storage media includes, but is not limited to, RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVDs) or other optical disk storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which can be used to store the desired information and which can be accessed by computing device 600. Computer storage media does not comprise signals per se. Communication media typically embodies computer-readable instructions, data structures, program modules, or other data in a modulated data signal such as a carrier wave or other transport mechanism and includes any information delivery media. The term “modulated data signal” means a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, communication media includes wired media, such as a wired network or direct-wired connection, and wireless media, such as acoustic, RF, infrared, and other

wireless media. Combinations of any of the above should also be included within the scope of computer-readable media.

[0139] Memory 612 includes computer storage media in the form of volatile and/or nonvolatile memory. The memory may be removable, non-removable, or a combination thereof. Exemplary hardware devices include solid-state memory, hard drives, optical-disc drives, etc. Computing device 600 includes one or more processors 614 that read data from various entities such as memory 612 or I/O components 620. Presentation component(s) 616 presents data indications to a user or other device. In some implementations, presentation component 220 of system 200 may be embodied as a presentation component 616. Other examples of presentation components may include a display device, speaker, printing component, vibrating component, and the like.

[0140] The I/O ports 618 allow computing device 600 to be logically coupled to other devices, including I/O components 620, some of which may be built in. Illustrative components include a microphone, joystick, game pad, satellite dish, scanner, printer, wireless device, etc. The I/O components 620 may provide a natural user interface (NUI) that processes air gestures, voice, or other physiological inputs generated by a user. In some instances, inputs may be transmitted to an appropriate network element for further processing. An NUI may implement any combination of speech recognition, touch and stylus recognition, facial recognition, biometric recognition, gesture recognition both on screen and adjacent to the screen, air gestures, head and eye tracking, and touch recognition associated with displays on the computing device 600. The computing device 600 may be equipped with depth cameras, such as stereoscopic camera systems, infrared camera systems, RGB camera systems, and combinations of these, for gesture detection and recognition. Additionally, the computing device 600 may be equipped with accelerometers or gyroscopes that enable detection of motion. The output of the accelerometers or gyroscopes may be provided to the display of the computing device 600 to render immersive augmented reality or virtual reality.

[0141] Some embodiments of computing device 600 may include one or more radio(s) 624 (or similar wireless communication components). The radio 624 transmits and receives radio or wireless communications. The computing device 600 may be a wireless terminal adapted to receive communications and media over various wireless networks. Computing device 600 may communicate via wireless protocols, such as code division multiple access (“CDMA”), global system for mobiles (“GSM”), or time division multiple access (“TDMA”), as well as others, to communicate with other devices. The radio communications may be a short-range connection, a long-range connection, or a combination of both a short-range and a long-range wireless telecommunications connection. When we refer to “short” and “long” types of connections, we do not mean to refer to the spatial relation between two devices. Instead, we are generally referring to short range and long range as different categories, or types, of connections (i.e., a primary connection and a secondary connection). A short-range connection may include, by way of example and not limitation, a Wi-Fi® connection to a device (e.g., mobile hotspot) that provides access to a wireless communications network, such as a WLAN connection using the 802.11 protocol; a Blu-

etooth connection to another computing device is a second example of a short-range connection, or a near-field communication connection. A long-range connection may include a connection using, by way of example and not limitation, one or more of CDMA, GPRS, GSM, TDMA, and 802.16 protocols.

[0142] Many different arrangements of the various components depicted, as well as components not shown, are possible without departing from the scope of the claims below. Embodiments of the disclosure have been described with the intent to be illustrative rather than restrictive. Alternative embodiments will become apparent to readers of this disclosure after and because of reading it. Alternative means of implementing the aforementioned can be completed without departing from the scope of the claims below. Certain features and sub-combinations are of utility and may be employed without reference to other features and sub-combinations and are contemplated within the scope of the claims.

[0143] Accordingly, in one aspect, an embodiment of the present disclosure is directed to a computerized system comprising one or more sensors configured to provide sensor data; a user activity monitor configured to identify and monitor a user device and user activity associated with the user device; an activity pattern inference engine configured to determine an activity pattern from a plurality of activity events; one or more processors; and computer storage memory having computer-executable instructions stored thereon which, when executed by the processor, implement a method of inferring a user activity pattern. The method includes (a) identifying a set of user devices associated with a user; (b) monitoring, using the user activity monitor, the set of user devices to detect a user activity event; (c) upon detecting a user activity event, determining a set of activity features associated with the activity event, the set of activity features determined based at least in part on the sensor data; (d) storing a record of the activity event and associated activity features in an activity event data store that comprises records of a plurality of activity events; (e) using the activity pattern inference engine to identify an activity pattern based on an analysis of the plurality of activity events to determine a set of activity events having similar activity features; and (f) storing a description of the activity pattern in an activity pattern data store.

[0144] In some embodiments of this system, the activity event comprises one or more of browsing to a website, launching an application, initiating a communication, performing a search query, setting a reminder, or scheduling an event on a calendar and the set of associated activity features comprise one or more features related to content associated with the activity event (a content feature), date or time of the activity event (date-time feature), location of the activity event (location feature), device-usage related features associated with the activity event (device-usage feature), or contextual information associated with the activity event. Further, in some embodiments, the activity features are determined to be similar based on a comparison of the same category of activity features using a similarity threshold that is predetermined according to the category of the compared activity features, the categories comprising, by way of example and not limitation, content features, date-time features, location features, device-usage features, or other types of activity features described herein.

[0145] In another aspect, an embodiment of the present disclosure is directed to a computerized system comprising one or more sensors configured to provide sensor data, each sensor associated with a user device in a set of user devices associated with a user; one or more processors; and computer storage memory having computer-executable instructions stored thereon which, when executed by the processor, implement a method of inferring a probable future user action. The method includes: (a) accessing an inferred user activity pattern for a user; (b) predicting a probable future activity event based on the activity pattern and a context determined at least in part from the sensor data; and (c) providing, using at least one user device in the set of user devices, an enhanced user experience based on the determined probable future activity event.

[0146] In one embodiment of this system, the sensor data for predicting a probable future activity event is provided from a sensor associated with a first user device in the set of user devices, and the enhanced user experience is provided using a second user device in the set of user devices. In one embodiment of this system, the provided enhanced user experience comprises one of a recommendation, notification, request, or suggestion related to the probable future activity. In another embodiment, the provided enhanced user experience comprises automatically carrying out the probable future activity at a time or location consistent with the inferred user activity pattern. In yet another embodiment, the predicted probable future activity event comprises connecting, at a future time, a user device having a battery to a power source to charge the battery, and the enhanced user experience comprises a battery power management service that reduces power consumption by the user device based on the future time.

[0147] In yet another aspect, an embodiment of the disclosure is directed to a computerized method of providing an enhanced user experience using one or more computing devices associated with a user. The method comprises identifying a set of computing devices associated with the user, at least a subset of the set of computing devices having one or more sensors configured to provide sensor data; and determining an activity pattern of user activity associated with a first computing device in the set of computing devices by: (a) monitoring, using a user activity monitor, the set of computing devices to detect a plurality of user activity events; (b) upon detecting each user activity event, determining a set of activity features associated with the activity event, the set of activity features determined based at least in part on sensor data from at least a first computing device in the subset of computing devices having one or more sensors; and (c) determining an activity pattern based on an analysis of the plurality of user activity events to determine a set of activity events having similar activity features. The method further comprises predicting a probable future activity event based on the determined activity pattern and a context determined at least in part from the sensor data; determining a confidence score associated with the predicted probable future activity event based at least in part on the number of user activity events used to determine the pattern, the confidence score indicating a probability that the future activity event will occur; and providing, using at least one computing device in the set of computing devices, an enhanced user experience based on the determined probable future activity event and corresponding confidence score. In some embodiments of this method, the enhanced user experience

rience is provided using a second computing device in the set of computing devices that is different than the first computing device.

What is claimed is:

1. At least one computer-storage medium having computer usable instructions embodied thereon which, when executed by at least one processor, implement a method comprising:

accessing an inferred user activity pattern for a user, the inferred user activity pattern associated with how a user interacts with a user device and comprising a device usage-related feature and a context feature, wherein the inferred user activity pattern comprises an indication of a time of day for charging the user device;

identifying both an amount of time remaining until the time of day for charging the user device and a mean consumption of battery energy;

predicting a probable future activity event to occur at a predicted future time, wherein the probable future activity event comprises that the user device will exhaust battery power prior to the time of day for charging the user device; and

providing, using the user device, an enhanced user experience based on the probable future activity event by causing presentation of a notification recommending that the user charge the user device.

2. The at least one computer-storage medium of claim 1, wherein causing presentation of the notification occurs responsive to a determination that the user is located at a location where the user has previously charged the user device.

3. At least one computer-storage medium having computer usable instructions embodied thereon which, when executed by at least one processor, implement a method comprising:

accessing an inferred user activity pattern for a user, the inferred user activity pattern associated with how a user interacts with a user device and comprising a device usage-related feature and a context feature, wherein the inferred user activity pattern comprises an indication of an activity performed on the user device manually by the user;

predicting a probable future activity event to occur at a predicted future time, wherein the probable future activity event comprises performance of the activity by the user; and

providing, using the user device, an enhanced user experience based on the probable future activity event by causing the activity to be performed automatically on the user device.

* * * * *