

# US Patent & Trademark Office

## Patent Public Search | Text View

---

United States Patent Application Publication

20250265507

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Betteridge; Jared Lamont et al.

---

### Machine Learning For Intent Matching Engine Comprising An Artificial Neural Network

---

#### Abstract

A server obtains a natural language query from a first user device. The server matches an intent to the natural language query using an intent matching engine. The intent represents predicted data associated with the natural language query. The server transmits the natural language query and the intent to a second user device. The server receives, from the second user device, a response indicating whether the natural language query is properly matched to the intent. The server trains the intent matching engine based on a machine learning technique and the response.

---

**Inventors:** Betteridge; Jared Lamont (Herriman, UT), Betteridge; Justin Bryce (Boise, ID), Brinton; Connor Isaac (Holly Springs, NC), Vong; Victor The (Dublin, CA), Wenke; Samuel John (Cincinnati, OH)

**Applicant:** Zoom Communications, Inc. (San Jose, CA)

**Family ID:** 1000008575217

**Appl. No.:** 19/198670

**Filed:** May 05, 2025

#### Related U.S. Application Data

parent US continuation 17969816 20221020 parent-grant-document US 12321830 child US 19198670

---

#### Publication Classification

**Int. Cl.:** G06N20/00 (20190101); G06F16/334 (20250101)

**U.S. Cl.:**

## Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION(S) [0001] This application is a continuation of U.S. patent application Ser. No. 17/969,816, filed Oct. 20, 2022, the entire disclosure of which is hereby incorporated by reference.

### FIELD

[0002] This disclosure relates to machine learning techniques for an intent matching engine, which may, for example, match a natural language query to an intent.

---

## Description

### BRIEF DESCRIPTION OF THE DRAWINGS

[0003] This disclosure is best understood from the following detailed description when read in conjunction with the accompanying drawings. It is emphasized that, according to common practice, the various features of the drawings are not to-scale. On the contrary, the dimensions of the various features are arbitrarily expanded or reduced for clarity.

[0004] FIG. 1 is a block diagram of an example of an electronic computing and communications system.

[0005] FIG. 2 is a block diagram of an example internal configuration of a computing device of an electronic computing and communications system.

[0006] FIG. 3 is a block diagram of an example of a software platform implemented by an electronic computing and communications system.

[0007] FIG. 4 is a block diagram of an example of a contact center system.

[0008] FIG. 5 is a block diagram of an example of a system in which intent matching may be performed.

[0009] FIG. 6 is a data flow diagram of an example of online learning for an intent matching engine.

[0010] FIG. 7 is a flow diagram of an example of processing client interface data.

[0011] FIG. 8 is a user interface diagram of an example of a query reviewing interface.

[0012] FIG. 9 is a flowchart of an example of a technique for online learning for an intent matching engine.

### DETAILED DESCRIPTION

[0013] A contact center may be operated by or otherwise on behalf of a business to assist users of various products and related services of the business. The contact center may be accessible to such users via multiple modalities, for example, telephone or Internet-based text, audio, or video communication implemented via a Unified Communication as a Service (UCaaS) system. In recent years, many contact centers have begun to use automated chatbots (e.g., computer-implemented engines that automatically process spoken or written natural language input provided by a user) that process queries in the form of audio or text-based input from a user in order to identify the user's goals and to provide an appropriate response. However, mapping a natural language audio or text-based input to an appropriate goal may be challenging. Techniques for automatically identifying the user's goals based on a natural language audio or text-based input may be desirable.

[0014] One solution to this problem is to search a data repository based on the natural language text (typed by the user or derived from an audio input) provided by the user. The data repository may be searched to respond to a query of the user and/or to meet a goal of the user (as indicated by the

query). However, this solution might not be effective as the natural language text might map to irrelevant contact (e.g., a query for “mobile phone” might lead to a result related to a mobile phone accessory) and/or might leave out relevant or otherwise related content (e.g., a query for “mobile phone” might not lead to a relevant result that describes a “cellular device”). Furthermore, if the query is in a first natural language (e.g., Spanish) and the data repository includes content in a second natural language (e.g., English), the natural language text-based search might not result in any results being generated.

[0015] Implementations of this disclosure address problems such as these by having a server access a natural language query from a querying user device. For example, a querying user might type, “Help me change my home address,” into a chat application to communicate with for a credit card company. The chat application may be accessible via a mobile application or website of the credit card company, and the communications may be processed by a chatbot. The server matches an intent to the natural language query using an intent matching engine. The intent may be “change billing address.” An intent may represent a goal of a user entering the natural language query, which could be handled by a specific workflow. Many different natural language queries may match to the same intent. For example, the spoken or written natural language queries, “Help me change my home address,” and “Update billing address,” may both map to the intent “change billing address.” The server maps a workflow to the intent. The workflow may be associated with a webpage, a set of human agents, or an automated workflow performed by the chatbot. The server transmits an output associated with the workflow to the querying user device. The output may include, for example, a webpage with instructions for how to change the billing address in the credit card company's online banking application. Alternatively, the output may include the chatbot asking the querying user to provide their new billing address and automatically updating the records of the credit card company. In another alternative, the output includes the user being transferred to a human agent who is able to assist with changing the billing address.

[0016] As used herein, the term “intent” may include a member of a set of stored intents, where multiple natural language queries may map to the same intent. The intent may correspond to predicted data, for example, a predicted goal, of a user entering the natural language query. The predicted data may be associated with the natural language query. For example, a credit card company may have intents including: “change billing address,” “dispute a charge,” “request credit limit increase,” and “request interest rate reduction,” each of which could have multiple different corresponding natural language queries.

[0017] The natural language query and the matched intent may be transmitted to a reviewing user device for review thereat. For example, the reviewing user device may be operated by a quality assurance employee of the credit card company. The user of the reviewing user device generates a reviewing response indicating whether the matching of the intent to the natural language query is accurate. The intent matching engine is trained using online learning based on the reviewing response. As a result of the online learning, the intent matching engine's performance may improve over time, and the intent matching engine might be able to adjust its performance responsive to changes in usage of the natural language over time.

[0018] According to some implementations, the server monitors activity of the querying user device with respect to the transmitted workflow. The intent matching engine is further trained based on the activity of the querying user device with respect to the transmitted workflow. For example, if the transmitted workflow is an article describing how to update the billing address, the server may monitor whether the user opened or scrolled through the article, and whether the user later logged into online banking to perform the steps indicated in the article.

[0019] In some cases, as multiple intents are identified from many queries of many users, a dashboard of intents over time may be generated. The dashboard may be transmitted to the reviewing user device or to an administrator device. The dashboard may indicate (e.g., in a graph or a table) a number of queries associated with each intent during multiple time periods. As a result,

the reviewing user or the administrator may be able to detect which intents are becoming more or less common with time. This may be a leading indicator of a trend. For example, if the number of users who have the intent “request interest rate reduction” is increasing, the credit card company may consider offering a lower interest rate product.

[0020] To describe some implementations in greater detail, reference is first made to examples of hardware and software structures used to implement online learning for an intent matching engine. FIG. 1 is a block diagram of an example of an electronic computing and communications system **100**, which can be or include a distributed computing system (e.g., a client-server computing system), a cloud computing system, a clustered computing system, or the like.

[0021] The system **100** includes one or more customers, such as customers **102A** through **102B**, which may each be a public entity, private entity, or another corporate entity or individual that purchases or otherwise uses software services, such as of a UCaaS platform provider. Each customer can include one or more clients. For example, as shown and without limitation, the customer **102A** can include clients **104A** through **104B**, and the customer **102B** can include clients **104C** through **104D**. A customer can include a customer network or domain. For example, and without limitation, the clients **104A** through **104B** can be associated or communicate with a customer network or domain for the customer **102A** and the clients **104C** through **104D** can be associated or communicate with a customer network or domain for the customer **102B**.

[0022] A client, such as one of the clients **104A** through **104D**, may be or otherwise refer to one or both of a client device or a client application. Where a client is or refers to a client device, the client can comprise a computing system, which can include one or more computing devices, such as a mobile phone, a tablet computer, a laptop computer, a notebook computer, a desktop computer, or another suitable computing device or combination of computing devices. Where a client instead is or refers to a client application, the client can be an instance of software running on a customer device (e.g., a client device or another device). In some implementations, a client can be implemented as a single physical unit or as a combination of physical units. In some implementations, a single physical unit can include multiple clients.

[0023] The system **100** can include a number of customers and/or clients or can have a configuration of customers or clients different from that generally illustrated in FIG. 1. For example, and without limitation, the system **100** can include hundreds or thousands of customers, and at least some of the customers can include or be associated with a number of clients.

[0024] The system **100** includes a datacenter **106**, which may include one or more servers. The datacenter **106** can represent a geographic location, which can include a facility, where the one or more servers are located. The system **100** can include a number of datacenters and servers or can include a configuration of datacenters and servers different from that generally illustrated in FIG. 1. For example, and without limitation, the system **100** can include tens of datacenters, and at least some of the datacenters can include hundreds or another suitable number of servers. In some implementations, the datacenter **106** can be associated or communicate with one or more datacenter networks or domains, which can include domains other than the customer domains for the customers **102A** through **102B**.

[0025] The datacenter **106** includes servers used for implementing software services of a UCaaS platform. The datacenter **106** as generally illustrated includes an application server **108**, a database server **110**, and a telephony server **112**. The servers **108** through **112** can each be a computing system, which can include one or more computing devices, such as a desktop computer, a server computer, or another computer capable of operating as a server, or a combination thereof. A suitable number of each of the servers **108** through **112** can be implemented at the datacenter **106**. The UCaaS platform uses a multi-tenant architecture in which installations or instantiations of the servers **108** through **112** is shared amongst the customers **102A** through **102B**.

[0026] In some implementations, one or more of the servers **108** through **112** can be a non-hardware server implemented on a physical device, such as a hardware server. In some

implementations, a combination of two or more of the application server **108**, the database server **110**, and the telephony server **112** can be implemented as a single hardware server or as a single non-hardware server implemented on a single hardware server. In some implementations, the datacenter **106** can include servers other than or in addition to the servers **108** through **112**, for example, a media server, a proxy server, or a web server.

[0027] The application server **108** runs web-based software services deliverable to a client, such as one of the clients **104A** through **104D**. As described above, the software services may be of a UCaaS platform. For example, the application server **108** can implement all or a portion of a UCaaS platform, including conferencing software, messaging software, and/or other intra-party or inter-party communications software. The application server **108** may, for example, be or include a unitary Java Virtual Machine (JVM).

[0028] In some implementations, the application server **108** can include an application node, which can be a process executed on the application server **108**. For example, and without limitation, the application node can be executed in order to deliver software services to a client, such as one of the clients **104A** through **104D**, as part of a software application. The application node can be implemented using processing threads, virtual machine instantiations, or other computing features of the application server **108**. In some such implementations, the application server **108** can include a suitable number of application nodes, depending upon a system load or other characteristics associated with the application server **108**. For example, and without limitation, the application server **108** can include two or more nodes forming a node cluster. In some such implementations, the application nodes implemented on a single application server **108** can run on different hardware servers.

[0029] The database server **110** stores, manages, or otherwise provides data for delivering software services of the application server **108** to a client, such as one of the clients **104A** through **104D**. In particular, the database server **110** may implement one or more databases, tables, or other information sources suitable for use with a software application implemented using the application server **108**. The database server **110** may include a data storage unit accessible by software executed on the application server **108**. A database implemented by the database server **110** may be a relational database management system (RDBMS), an object database, an XML database, a configuration management database (CMDB), a management information base (MIB), one or more flat files, other suitable non-transient storage mechanisms, or a combination thereof. The system **100** can include one or more database servers, in which each database server can include one, two, three, or another suitable number of databases configured as or comprising a suitable database type or combination thereof.

[0030] In some implementations, one or more databases, tables, other suitable information sources, or portions or combinations thereof may be stored, managed, or otherwise provided by one or more of the elements of the system **100** other than the database server **110**, for example, the client **104** or the application server **108**.

[0031] The telephony server **112** enables network-based telephony and web communications from and to clients of a customer, such as the clients **104A** through **104B** for the customer **102A** or the clients **104C** through **104D** for the customer **102B**. Some or all of the clients **104A** through **104D** may be voice over internet protocol (VOIP)-enabled devices configured to send and receive calls over a network **114**. In particular, the telephony server **112** includes a session initiation protocol (SIP) zone and a web zone. The SIP zone enables a client of a customer, such as the customer **102A** or **102B**, to send and receive calls over the network **114** using SIP requests and responses. The web zone integrates telephony data with the application server **108** to enable telephony-based traffic access to software services run by the application server **108**. Given the combined functionality of the SIP zone and the web zone, the telephony server **112** may be or include a cloud-based private branch exchange (PBX) system.

[0032] The SIP zone receives telephony traffic from a client of a customer and directs same to a

destination device. The SIP zone may include one or more call switches for routing the telephony traffic. For example, to route a VOIP call from a first VOIP-enabled client of a customer to a second VOIP-enabled client of the same customer, the telephony server **112** may initiate a SIP transaction between a first client and the second client using a PBX for the customer. However, in another example, to route a VOIP call from a VOIP-enabled client of a customer to a client or non-client device (e.g., a desktop phone which is not configured for VOIP communication) which is not VOIP-enabled, the telephony server **112** may initiate a SIP transaction via a VOIP gateway that transmits the SIP signal to a public switched telephone network (PSTN) system for outbound communication to the non-VOIP-enabled client or non-client phone. Hence, the telephony server **112** may include a PSTN system and may in some cases access an external PSTN system.

[0033] The telephony server **112** includes one or more session border controllers (SBCs) for interfacing the SIP zone with one or more aspects external to the telephony server **112**. In particular, an SBC can act as an intermediary to transmit and receive SIP requests and responses between clients or non-client devices of a given customer with clients or non-client devices external to that customer. When incoming telephony traffic for delivery to a client of a customer, such as one of the clients **104A** through **104D**, originating from outside the telephony server **112** is received, a SBC receives the traffic and forwards it to a call switch for routing to the client.

[0034] In some implementations, the telephony server **112**, via the SIP zone, may enable one or more forms of peering to a carrier or customer premise. For example, Internet peering to a customer premise may be enabled to ease the migration of the customer from a legacy provider to a service provider operating the telephony server **112**. In another example, private peering to a customer premise may be enabled to leverage a private connection terminating at one end at the telephony server **112** and at the other end at a computing aspect of the customer environment. In yet another example, carrier peering may be enabled to leverage a connection of a peered carrier to the telephony server **112**.

[0035] In some such implementations, a SBC or telephony gateway within the customer environment may operate as an intermediary between the SBC of the telephony server **112** and a PSTN for a peered carrier. When an external SBC is first registered with the telephony server **112**, a call from a client can be routed through the SBC to a load balancer of the SIP zone, which directs the traffic to a call switch of the telephony server **112**. Thereafter, the SBC may be configured to communicate directly with the call switch.

[0036] The web zone receives telephony traffic from a client of a customer, via the SIP zone, and directs same to the application server **108** via one or more Domain Name System (DNS) resolutions. For example, a first DNS within the web zone may process a request received via the SIP zone and then deliver the processed request to a web service which connects to a second DNS at or otherwise associated with the application server **108**. Once the second DNS resolves the request, it is delivered to the destination service at the application server **108**. The web zone may also include a database for authenticating access to a software application for telephony traffic processed within the SIP zone, for example, a softphone.

[0037] The clients **104A** through **104D** communicate with the servers **108** through **112** of the datacenter **106** via the network **114**. The network **114** can be or include, for example, the Internet, a local area network (LAN), a wide area network (WAN), a virtual private network (VPN), or another public or private means of electronic computer communication capable of transferring data between a client and one or more servers. In some implementations, a client can connect to the network **114** via a communal connection point, link, or path, or using a distinct connection point, link, or path. For example, a connection point, link, or path can be wired, wireless, use other communications technologies, or a combination thereof.

[0038] The network **114**, the datacenter **106**, or another element, or combination of elements, of the system **100** can include network hardware such as routers, switches, other network devices, or combinations thereof. For example, the datacenter **106** can include a load balancer **116** for routing

traffic from the network **114** to various servers associated with the datacenter **106**. The load balancer **116** can route, or direct, computing communications traffic, such as signals or messages, to respective elements of the datacenter **106**.

[0039] For example, the load balancer **116** can operate as a proxy, or reverse proxy, for a service, such as a service provided to one or more remote clients, such as one or more of the clients **104A** through **104D**, by the application server **108**, the telephony server **112**, and/or another server. Routing functions of the load balancer **116** can be configured directly or via a DNS. The load balancer **116** can coordinate requests from remote clients and can simplify client access by masking the internal configuration of the datacenter **106** from the remote clients.

[0040] In some implementations, the load balancer **116** can operate as a firewall, allowing or preventing communications based on configuration settings. Although the load balancer **116** is depicted in FIG. **1** as being within the datacenter **106**, in some implementations, the load balancer **116** can instead be located outside of the datacenter **106**, for example, when providing global routing for multiple datacenters. In some implementations, load balancers can be included both within and outside of the datacenter **106**. In some implementations, the load balancer **116** can be omitted.

[0041] FIG. **2** is a block diagram of an example internal configuration of a computing device **200** of an electronic computing and communications system. In one configuration, the computing device **200** may implement one or more of the client **104**, the application server **108**, the database server **110**, or the telephony server **112** of the system **100** shown in FIG. **1**.

[0042] The computing device **200** includes components or units, such as a processor **202**, a memory **204**, a bus **206**, a power source **208**, peripherals **210**, a user interface **212**, a network interface **214**, other suitable components, or a combination thereof. One or more of the memory **204**, the power source **208**, the peripherals **210**, the user interface **212**, or the network interface **214** can communicate with the processor **202** via the bus **206**.

[0043] The processor **202** is a central processing unit, such as a microprocessor, and can include single or multiple processors having single or multiple processing cores. Alternatively, the processor **202** can include another type of device, or multiple devices, configured for manipulating or processing information. For example, the processor **202** can include multiple processors interconnected in one or more manners, including hardwired or networked. The operations of the processor **202** can be distributed across multiple devices or units that can be coupled directly or across a local area or other suitable type of network. The processor **202** can include a cache, or cache memory, for local storage of operating data or instructions.

[0044] The memory **204** includes one or more memory components, which may each be volatile memory or non-volatile memory. For example, the volatile memory can be random access memory (RAM) (e.g., a DRAM module, such as DDR SDRAM). In another example, the non-volatile memory of the memory **204** can be a disk drive, a solid state drive, flash memory, or phase-change memory. In some implementations, the memory **204** can be distributed across multiple devices. For example, the memory **204** can include network-based memory or memory in multiple clients or servers performing the operations of those multiple devices.

[0045] The memory **204** can include data for immediate access by the processor **202**. For example, the memory **204** can include executable instructions **216**, application data **218**, and an operating system **220**. The executable instructions **216** can include one or more application programs, which can be loaded or copied, in whole or in part, from non-volatile memory to volatile memory to be executed by the processor **202**. For example, the executable instructions **216** can include instructions for performing some or all of the techniques of this disclosure. The application data **218** can include user data, database data (e.g., database catalogs or dictionaries), or the like. In some implementations, the application data **218** can include functional programs, such as a web browser, a web server, a database server, another program, or a combination thereof. The operating system **220** can be, for example, Microsoft Windows®, Mac OS X®, or Linux®; an operating

system for a mobile device, such as a smartphone or tablet device; or an operating system for a non-mobile device, such as a mainframe computer.

[0046] The power source **208** provides power to the computing device **200**. For example, the power source **208** can be an interface to an external power distribution system. In another example, the power source **208** can be a battery, such as where the computing device **200** is a mobile device or is otherwise configured to operate independently of an external power distribution system. In some implementations, the computing device **200** may include or otherwise use multiple power sources. In some such implementations, the power source **208** can be a backup battery.

[0047] The peripherals **210** includes one or more sensors, detectors, or other devices configured for monitoring the computing device **200** or the environment around the computing device **200**. For example, the peripherals **210** can include a geolocation component, such as a global positioning system location unit. In another example, the peripherals can include a temperature sensor for measuring temperatures of components of the computing device **200**, such as the processor **202**. In some implementations, the computing device **200** can omit the peripherals **210**.

[0048] The user interface **212** includes one or more input interfaces and/or output interfaces. An input interface may, for example, be a positional input device, such as a mouse, touchpad, touchscreen, or the like; a keyboard; or another suitable human or machine interface device. An output interface may, for example, be a display, such as a liquid crystal display, a cathode-ray tube, a light emitting diode display, or other suitable display.

[0049] The network interface **214** provides a connection or link to a network (e.g., the network **114** shown in FIG. **1**). The network interface **214** can be a wired network interface or a wireless network interface. The computing device **200** can communicate with other devices via the network interface **214** using one or more network protocols, such as using Ethernet, transmission control protocol (TCP), internet protocol (IP), power line communication, an IEEE 802.X protocol (e.g., Wi-Fi, Bluetooth, or ZigBee), infrared, visible light, general packet radio service (GPRS), global system for mobile communications (GSM), code-division multiple access (CDMA), Z-Wave, another protocol, or a combination thereof.

[0050] FIG. **3** is a block diagram of an example of a software platform **300** implemented by an electronic computing and communications system, for example, the system **100** shown in FIG. **1**. The software platform **300** is a UCaaS platform accessible by clients of a customer of a UCaaS platform provider, for example, the clients **104A** through **104B** of the customer **102A** or the clients **104C** through **104D** of the customer **102B** shown in FIG. **1**. The software platform **300** may be a multi-tenant platform instantiated using one or more servers at one or more datacenters including, for example, the application server **108**, the database server **110**, and the telephony server **112** of the datacenter **106** shown in FIG. **1**.

[0051] The software platform **300** includes software services accessible using one or more clients. For example, a customer **302** as shown includes four clients—a desk phone **304**, a computer **306**, a mobile device **308**, and a shared device **310**. The desk phone **304** is a desktop unit configured to at least send and receive calls and includes an input device for receiving a telephone number or extension to dial to and an output device for outputting audio and/or video for a call in progress. The computer **306** is a desktop, laptop, or tablet computer including an input device for receiving some form of user input and an output device for outputting information in an audio and/or visual format. The mobile device **308** is a smartphone, wearable device, or other mobile computing aspect including an input device for receiving some form of user input and an output device for outputting information in an audio and/or visual format. The desk phone **304**, the computer **306**, and the mobile device **308** may generally be considered personal devices configured for use by a single user. The shared device **310** is a desk phone, a computer, a mobile device, or a different device which may instead be configured for use by multiple specified or unspecified users.

[0052] Each of the clients **304** through **310** includes or runs on a computing device configured to access at least a portion of the software platform **300**. In some implementations, the customer **302**



may include additional clients not shown. For example, the customer **302** may include multiple clients of one or more client types (e.g., multiple desk phones or multiple computers) and/or one or more clients of a client type not shown in FIG. 3 (e.g., wearable devices or televisions other than as shared devices). For example, the customer **302** may have tens or hundreds of desk phones, computers, mobile devices, and/or shared devices.

[0053] The software services of the software platform **300** generally relate to communications tools, but are in no way limited in scope. As shown, the software services of the software platform **300** include telephony software **312**, conferencing software **314**, messaging software **316**, and other software **318**. Some or all of the software **312** through **318** uses customer configurations **320** specific to the customer **302**. The customer configurations **320** may, for example, be data stored within a database or other data store at a database server, such as the database server **110** shown in FIG. 1.

[0054] The telephony software **312** enables telephony traffic between ones of the clients **304** through **310** and other telephony-enabled devices, which may be other ones of the clients **304** through **310**, other VOIP-enabled clients of the customer **302**, non-VOIP-enabled devices of the customer **302**, VOIP-enabled clients of another customer, non-VOIP-enabled devices of another customer, or other VOIP-enabled clients or non-VOIP-enabled devices. Calls sent or received using the telephony software **312** may, for example, be sent or received using the desk phone **304**, a softphone running on the computer **306**, a mobile application running on the mobile device **308**, or using the shared device **310** that includes telephony features.

[0055] The telephony software **312** further enables phones that do not include a client application to connect to other software services of the software platform **300**. For example, the telephony software **312** may receive and process calls from phones not associated with the customer **302** to route that telephony traffic to one or more of the conferencing software **314**, the messaging software **316**, or the other software **318**.

[0056] The conferencing software **314** enables audio, video, and/or other forms of conferences between multiple participants, such as to facilitate a conference between those participants. In some cases, the participants may all be physically present within a single location, for example, a conference room, in which the conferencing software **314** may facilitate a conference between only those participants and using one or more clients within the conference room. In some cases, one or more participants may be physically present within a single location and one or more other participants may be remote, in which the conferencing software **314** may facilitate a conference between all of those participants using one or more clients within the conference room and one or more remote clients. In some cases, the participants may all be remote, in which the conferencing software **314** may facilitate a conference between the participants using different clients for the participants. The conferencing software **314** can include functionality for hosting, presenting scheduling, joining, or otherwise participating in a conference. The conferencing software **314** may further include functionality for recording some or all of a conference and/or documenting a transcript for the conference.

[0057] The messaging software **316** enables instant messaging, unified messaging, and other types of messaging communications between multiple devices, such as to facilitate a chat or other virtual conversation between users of those devices. The unified messaging functionality of the messaging software **316** may, for example, refer to email messaging which includes a voicemail transcription service delivered in email format.

[0058] The other software **318** enables other functionality of the software platform **300**. Examples of the other software **318** include, but are not limited to, device management software, resource provisioning and deployment software, administrative software, third party integration software, and the like. In one particular example, the other software **318** can include software of the intent management engine and software for online learning by the intent management engine.

[0059] The software **312** through **318** may be implemented using one or more servers, for example,

of a datacenter such as the datacenter **106** shown in FIG. **1**. For example, one or more of the software **312** through **318** may be implemented using an application server, a database server, and/or a telephony server, such as the servers **108** through **112** shown in FIG. **1**. In another example, one or more of the software **312** through **318** may be implemented using servers not shown in FIG. **1**, for example, a meeting server, a web server, or another server. In yet another example, one or more of the software **312** through **318** may be implemented using one or more of the servers **108** through **112** and one or more other servers. The software **312** through **318** may be implemented by different servers or by the same server.

[0060] Features of the software services of the software platform **300** may be integrated with one another to provide a unified experience for users. For example, the messaging software **316** may include a user interface element configured to initiate a call with another user of the customer **302**. In another example, the telephony software **312** may include functionality for elevating a telephone call to a conference. In yet another example, the conferencing software **314** may include functionality for sending and receiving instant messages between participants and/or other users of the customer **302**. In yet another example, the conferencing software **314** may include functionality for file sharing between participants and/or other users of the customer **302**. In some implementations, some or all of the software **312** through **318** may be combined into a single software application run on clients of the customer, such as one or more of the clients **304** through **310**.

[0061] FIG. **4** is a block diagram of an example of a contact center system. A contact center **400**, which in some cases may be implemented in connection with a software platform (e.g., the software platform **300** shown in FIG. **3**), is accessed by a user device **402** and used to establish a connection between the user device **402** and an agent device **404** over one of multiple modalities available for use with the contact center **400**, for example, telephony, video, text messaging, chat, and social media. The contact center **400** is implemented using one or more servers and software running thereon. For example, the contact center **400** may be implemented using one or more of the servers **108** through **112** shown in FIG. **1**, and may use communication software such as or similar to the software **312** through **318** shown in FIG. **3**. The contact center **400** includes software for facilitating contact center engagements requested by user devices such as the user device **402**. As shown, the software includes request processing software **406**, agent selection software **408**, and session handling software **410**.

[0062] The request processing software **406** processes a request for a contact center engagement initiated by the user device **402** to determine information associated with the request. The request may include a natural language query or a request entered in another manner (e.g., “press 1 to pay a bill, press 2 to request service”). The information associated with the request generally includes information identifying the purpose of the request and which is usable to direct the request traffic to a contact center agent capable of addressing the request. The information associated with the request may include information obtained from a user of the user device **402** after the request is initiated. For example, for the telephony modality, the request processing software **406** may use an interactive voice response (IVR) menu to prompt the user of the user device to present information associated with the purpose of the request, such as by identifying a category or sub-category of support requested. In another example, for the video modality, the request processing software **406** may use a form or other interactive user interface to prompt a user of the user device **402** to select options which correspond to the purpose of the request. In yet another example, for the chat modality, the request processing software **406** may ask the user of the user device **402** to summarize the purpose of the request (e.g., the natural language query) via text and thereafter process the text entered by the user device **402** using natural language processing and/or other processing.

[0063] The session handling software **410** establishes a connection between the user device **402** and the agent device **404**, which is the device of the agent selected by the agent selection software

**408.** The particular manner of the connection and the process for establishing same may be based on the modality used for the contact center engagement requested by the user device **402**. The contact center engagement is then facilitated over the established connection. For example, facilitating the contact center engagement over the established connection can include enabling the user of the user device **402** and the selected agent associated with the agent device **404** to engage in a discussion over the subject modality to address the purpose of the request from the user device **402**. The facilitation of the contact center engagement over the established connection can use communication software implemented in connection with a software platform, for example, one of the software **312** through **318**, or like software.

[0064] The user device **402** is a device configured to initiate a request for a contact center engagement which may be obtained and processed using the request processing software **406**. In some cases, the user device **402** may be a client device, for example, one of the clients **304** through **310** shown in FIG. 3. For example, the user device **402** may use a client application running thereat to initiate the request for the contact center engagement. In another example, the connection between the user device **402** and the agent device **404** may be established using software available to a client application running at the user device **402**. Alternatively, in some cases, the user device **402** may be other than a client device.

[0065] The agent device **404** is a device configured for use by a contact center agent. Where the contact center agent is a human, the agent device **404** is a device having a user interface. In some such cases, the agent device **404** may be a client device, for example, one of the clients **304** through **310**, or a non-client device. In some such cases, the agent device **404** may be a server which implements software usable by one or more contact center agents to address contact center engagements requested by contact center users. Where the contact center agent is a non-human, the agent device **404** is a device that may or may not have a user interface. For example, in some such cases, the agent device **404** may be a server which implements software of or otherwise usable in connection with the contact center **400**.

[0066] Although the request processing software **406**, the agent selection software **408**, and the session handling software **410** are shown as separate software components, in some implementations, some or all of the request processing software **406**, the agent selection software **408**, and the session handling software **410** may be combined. For example, the contact center **400** may be or include a single software component which performs the functionality of all of the request processing software **406**, the agent selection software **408**, and the session handling software **410**. In some implementations, one or more of the request processing software **406**, the agent selection software **408**, or the session handling software **410** may be comprised of multiple software components. In some implementations, the contact center **400** may include software components other than the request processing software **406**, the agent selection software **408**, and the session handling software **410**, such as in addition to or in place of one or more of the request processing software **406**, the agent selection software **408**, and the session handling software **410**.

[0067] FIG. 5 is a block diagram of an example of a system **500** in which intent matching may be performed. As shown, the system **500** includes a querying user device **502**, a server **504**, and a reviewing user device **506** capable of communicating over a network **508**. The server **504** stores an intent matching engine **510** that matches an input natural language query to an intent. The intent matching engine **510** may be implemented using machine learning or other artificial intelligence technologies. In some examples, the intent matching engine **510** leverages a convolutional neural network (CNN) that is capable of online learning. Alternatively, any other type of artificial neural network (ANN) may be used. In some cases, the intent matching engine may be implemented using at least one of a classification model, a regression model, clustering, dimensionality reduction, or deep learning. More details of the operation of the intent matching engine are provided below. As used herein, an engine may include software, hardware, or a combination of software and hardware.

[0068] The querying user device **502** may correspond to the user device **402** and/or to one of the clients **104A-D**. The server **504** may correspond to a server of the contact center **400** and/or to the application server **108**. The reviewing user device **506** may correspond to the agent device **404** and/or to one of the clients **104A-D**.

[0069] According to some implementations, the querying user device **502** transmits a natural language query to the server **504** over the network **508**. The natural language query may be spoken, typed, or otherwise inputted into the querying user device **502** text or recording in a natural language (e.g., English, Spanish, or Chinese). The natural language query may include a question or a request provided to a chatbot, for example, "I would like to cancel my subscription."

[0070] In response to receipt of the query, the server **504** uses the intent matching engine **510** to match the query to an intent, and transmits, to the querying user device **502**, an output associated with the intent. An intent may represent a goal or a purpose that a user of the querying user device is attempting to accomplish. For example, the query, "I would like to cancel my subscription," may match to the intent, "cancel subscription." The server **504** transmits (e.g., after the output has been transmitted to the querying user device **502**), to the reviewing user device **506**, an indication of the query received from the querying user device **502** and the intent matched to the query. The server **504** receives, from the reviewing user device **506**, a response indicating whether the query was properly matched to the intent. The server **504** further trains the intent matching engine **510**, using online learning, based on the response received from the reviewing user device **506**. The training may improve future operation of the intent matching engine **510**, but might not adjust the output that was provided to the querying user device **502** in response to the natural language query. Examples of operation of the querying user device **502**, the server **504**, and the receiving user device **506** are described in greater detail in conjunction with FIG. 6.

[0071] FIG. 6 is a data flow diagram of an example of online learning **600** for an intent matching engine. As shown, the online learning **600** is implemented using the querying user device **502**, the server **504**, and the reviewing user device **506** described in conjunction with FIG. 5.

[0072] As shown in FIG. 6, the querying user device **502** generates a query **602** in a natural language and transmits the query **602** to the server **504**. The query **602** may be generated during a communication with a chatbot of the server **504** via a telephone call, a messaging application, a website, or a dedicated application of an entity associated with the chatbot. For example, a user might wish to open a savings account at a bank. To accomplish this, the user might telephone the bank and may be asked, by a chatbot, "Please state the reason for your call." In response, the user might say, "I want to open a new savings account," and that phrase may become the query.

[0073] The server **504** matches the query **602** to an intent **604** using the intent matching engine **510** shown in FIG. 5. The intent **604** may be identified by applying a CNN or other artificial intelligence software of the intent matching engine **510** to the query **602**. For example, the query **602** "I want to open a new savings account," may match to the intent "open account." The server **504** identifies a workflow **606** corresponding to the intent **604**. The workflow **606** may include, for example, a link to a webpage for opening new accounts with the bank or connection with a human representative who can assist the user of the querying user device **502** with opening the new account. In some cases, the workflow **606** may be performed by the chatbot of the server **504** directly. For example, to open the new account, the workflow **606** may include the chatbot verifying the user's identity (e.g., by having the user provide their social security number and answer security questions generated based on their credit report), verifying the type of account (e.g., savings account) that the user wishes to open, requesting the account number of another account from which the opening deposit is to be obtained, and initiating an automated clearing house (ACH) request for the opening deposit. The server transmits a workflow output **608** associated with the workflow (e.g., questions by the chatbot, the link to the webpage, or the connection with the human representative) to the querying user device **502**.

[0074] The server **504** transmits a query-intent combination **610**, including the query **602** and the

intent **604**, to a reviewing user device **506**. The reviewing user device **506** may be operated by an agent associated with the server **504**, for example, an employee of a business (e.g., the bank) associated with the server. While the query-intent combination **610** is illustrated, it should be noted that the query **602** and the intent **604** may be transmitted to the reviewing user device **506** separately and may be stored separately at the reviewing user device **506**. The reviewing user device **506** may generate a visual output (e.g., as shown in FIG. **8**) representing the query-intent combination **610** and prompt the user of the reviewing user device **506** to input a reviewing response **612** indicating whether the query was accurately matched to the intent. For example, the query **602** “The balance on my checking account is high; I want to open a new savings account,” is accurately matched to the intent **604** “open account.” However, had this query **602** been matched to the intent “check account balance,” the user might indicate that the query was not accurately matched to the intent.

[0075] The reviewing user device **506** transmits the reviewing response **612** to the server **504**. At the server **504**, the reviewing response **612** is incorporated into training data **614**, which is used to further train the intent matching engine **510** to improve its performance in the future. The server **504** may use online learning techniques to recursively train the intent matching engine **510** based on the training data **614** which is generated from the operation of the intent matching engine **510**. As a result, if the intent matching engine **510** has little training data when it is initialized at the server **504**, the intent matching engine **510** may initially perform relatively poorly. However, the performance of the intent matching engine **510** may improve over time as more training data becomes available.

[0076] FIG. **7** is a flow diagram of an example of processing client interface data **700**. As shown, a client interface **702**, which includes data that may be transmitted to the clients **104A-D** or the querying user device **502** includes an announcement **704**, a suggestion **706**, and a query **708**. For example, in the context of an online store, an announcement may be transmitted (e.g., to subscribers who subscribe to receive email or messages (e.g., via short message service (SMS) or multimedia messaging service (MMS) through a cellular provider or via an online messaging service)). The announcement may include, for example, information about discounts or products available at the online store. The client interface may include a suggestion **706**. The suggestion may represent a recommendation by an employee of the online store or via an artificial intelligence engine operated by the online store. For example, a user who recently purchased a printer from the online store may receive a recommendation to purchase ink or paper for the printer.

Announcements **704** and suggestions **706** may be transmitted to users who requested to receive such notifications, and users may be able to opt out or unsubscribe from future announcements **704** or suggestions **706**. The client interface **702** may also include an interface for a user-generated query **708**. The query **708** may include a spoken or written natural language request that is automatically processed by a server. The query **708** may correspond to the query **602**.

[0077] As shown, the announcement **704** may include no user action or link to external site **710**. The suggestion **706** may link to a workflow **712** (e.g., the suggestion **706** to purchase an item from an online store may link directly to a workflow (e.g., a webpage or chatbot interface) for purchasing the item from the online store). Alternatively, the suggestion **706** may link to the creation of a query from the suggestion **714**. For example, a suggestion to purchase handyperson services to assist with assembling furniture that the user recently purchased from the online store for delivery to San Francisco may lead to the generation of a natural language query for finding a handyperson who is skilled in furniture assembly and located in San Francisco, with some parts of the query (e.g., a geographic location or a maximum travel distance) being input by the user.

[0078] Upon receiving the query **708** via the client interface **702**, the server **504** determines whether an intent is matched **716** to the query. If not, a non-intent based search **718** of the query **708** is conducted. For example, in the online store context, if the query includes the word “coffee,” and no intent is matched to “coffee,” the online store may be searched for products (or other pages)

matching “coffee.” If an intent is matched, the intent **720** is (at least temporarily) stored, and a workflow **722** corresponding to the intent **720** is identified. (It should be noted that the workflow may also be identified based on the suggestion **706** in the link to workflow **712**.) The intent **720** may correspond to the intent **604** of FIG. **6**. The workflow **722** may correspond to the workflow **606** of FIG. **6**.

[0079] The workflow **722** may include at least one of a support channel **724**, custom actions **726**, or a solution article. For example, in the online store context, the workflow **722** for the intent **720** “apply coupon,” may include the support channel **724** associated with a human agent who is capable of applying coupons to orders. Alternatively, the workflow **722** for the intent “apply coupon” may include custom actions **726**, implemented by a chatbot, to apply the coupon (e.g., asking the user to input the product they wish to purchase and the coupon number, and verifying whether the coupon was successfully applied). In another example, the workflow **722** for the intent “apply coupon” may include the solution article **728** which explains how to apply coupons via the mobile application or website of the online store.

[0080] FIG. **8** is a user interface diagram of an example of a query reviewing interface **800**. As shown, the query reviewing interface **800** is presented at a tablet computer **802**. In alternative implementations, the query reviewing interface **800** may be presented at a laptop computer, a desktop computer, a mobile phone or any other computing device. The tablet computer **802** displays a page **804** associated with the intent “process refund.” The page **804** includes queries for review **806**, which were received, at the server **504**, from querying user devices **502**, and matched, by the intent matching engine **510**, to the intent “process refund.” The queries for review **806** include queries **808A**, **808B**. Each query **808A**, **808B** is associated with a button **810A**, **810B** for indicating that the query **808A**, **808B** represents the intent “process refund” of the page **804**, and a button **812A**, **812B** for indicating that the query does not represent the intent “process refund” of the page **804**. As illustrated, the query **808A**, “I don't like the product, would it be possible to get a refund?” is properly matched to the intent “process refund,” so the user of the tablet computer **802** may select the “yes” button **810A**. The query **808B**, “Can I get a replacement instead of a refund?” is not properly matched to the intent “process refund.” Thus, the user of the tablet computer may select the “no” button **812B**.

[0081] While two queries are illustrated, the disclosed technique may be implemented with any number of queries, not necessarily two. As shown, both queries **808A**, **808B** are displayed in text. In alternative implementations, one or both of the queries **808A**, **808B** may be audio or video recordings. As show, both queries **808A**, **808B** are in English (and the intents and workflows are also in English). In alternative implementations, one or both of the queries may be in a different language (e.g., Spanish). For example, the text of the query **808A** may be, “No me gusta el producto, ¿sería posible obtener un reembolso?” which can be translated to “I don't like the product, would it be possible to get a refund?” in Spanish, and the intent matching engine **510** may still properly match the query **808A** to the English language intent “process refund.” In some implementations, the server **504** translates the query from Spanish into English prior to matching the intent to the query. In some implementations, the intent matching engine **510** is trained to process Spanish language queries, and does not translate the Spanish language query prior to matching it to the English language intent.

[0082] The workflow may associated with the intent may be available in English only, and the English language workflow may be provided in response to the Spanish language query. Alternatively, there may be multiple workflows in different languages, and a Spanish language workflow may be provided in response to the Spanish language query. In some cases, a Spanish language workflow might not be available, but workflows might be available in English and French. The user of the querying user device **502** may be prompted to indicate whether they prefer to proceed with an English or French workflow. The prompt may be presented in multiple different languages simultaneously, in Spanish, or in a default language (e.g., English) associated with the

server **504**.

[0083] To further describe some implementations in greater detail, reference is next made to examples of techniques which may be performed by or using an intent matching engine. FIG. **9** is a flowchart of an example of a technique **900** for online learning for an intent matching engine. The technique **900** can be executed using computing devices, such as the systems, hardware, and software described with respect to FIGS. **1-8**. The technique **900** can be performed, for example, by executing a machine-readable program or other computer-executable instructions, such as routines, instructions, programs, or other code. The steps, or operations, of the technique **900** or another technique, method, process, or algorithm described in connection with the implementations disclosed herein can be implemented directly in hardware, firmware, software executed by hardware, circuitry, or a combination thereof.

[0084] For simplicity of explanation, the technique **900** is depicted and described herein as a series of steps or operations. However, the steps or operations in accordance with this disclosure can occur in various orders and/or concurrently. Additionally, other steps or operations not presented and described herein may be used. Furthermore, not all illustrated steps or operations may be required to implement a technique in accordance with the disclosed subject matter.

[0085] At **902**, a server (e.g., the server **504**) obtains a natural language query from a first user device (e.g., the querying user device **502**). For example, a user of the first user device may speak or type the query to a chatbot during a communication session with the chatbot. The first user device may access the chatbot by connecting to the server using at least one of: a PSTN connection, a chat application, a dedicated application associated with the server, or a website associated with the server. The PSTN connection may correspond to a toll-free number of a business associated with the server (e.g., a toll free number of ABC Corporation). The chat application may include an account for the business (e.g., an account of ABC Corporation in a social media or messaging service). The dedicated application may be a mobile or tablet computer application of the business (e.g., a mobile application of ABC Corporation). The website may be a website of the business (e.g., a website of ABC Corporation).

[0086] At **904**, the server matches an intent to the natural language query using an intent matching engine (e.g., the intent matching engine **510**). The intent represents a predicted user goal associated with the natural language query. The server may store a set of intents, and may attempt to match the natural language query to one of the intents, or to determine that the natural language query does not match to any of the stored intents. If the natural language query does not match to any of the intents, other search techniques (e.g., text-based search) may be applied, an error message may be presented, or the natural language query may be forwarded to a human agent for manual processing. The intent matching engine may be an artificial neural network that is trained to match queries to intents based on text or utterances in the query. Alternatively, other artificial intelligence or rule-based technology may be used to implement the intent matching engine.

[0087] In some cases, the server stores a first set of global intents, that are applicable to multiple different businesses or organizations, and a second set of custom intents that are applicable to the business or organization operating the server. For example, the intent “change billing address” may be associated with multiple different businesses, for example, banks, insurance companies, online stores, and airlines. However, the intent “change seat” may be relevant to an airline but not relevant to a bank, an insurance company, or an online store. Thus, the set of intents for an airline may include global intents, such as, “change billing address,” and “reset password,” and some airline-specific intents, such as “change seat,” and “report lost luggage.”

[0088] The server may access a data repository that maps intents to workflows. In some implementations, the server transmits, to the first user device in response to the natural language query, a workflow output associated with a workflow to which the intent is mapped. The workflow output may include at least one of: a hyperlink, a webpage, an article, a list of actions, a chatbot-operated interface, or initiating a communication session between a human agent assigned to

handle the workflow and a user of the first user device. In some cases, an intent may map to a single workflow, and the output associated with the single workflow may be provided whenever the intent is encountered.

[0089] In some cases, an intent may map to multiple workflows. If the intent maps to multiple workflows, artificial intelligence (or other) techniques may be used to select a workflow for the user. For example, if the intent is “file a claim” for a car insurance business, and the user associated with the intent has previously used the website to manage their insurance policy, an article explaining how to file a claim may be provided to the user. Alternatively, if the user is a person who rarely accesses the website (or does not have a user account associated with the website), a human agent (e.g., operating the agent device **404**) may be connected with the user. In some examples, if the user is known to be fluent in German and to have difficulty communicating in English, a German version of the article and/or a German-speaking human agent may be used.

[0090] The server may monitor activity of the first user device with respect to the transmitted workflow output. The intent matching engine may be further trained based on activity of the first user device with respect to the transmitted workflow output. For example, if the first user device is provided with the hyperlink to the article about how to file a claim on the insurance company website, but the first user device never selects the article and the first user device is not used to file the claim, the server may determine that the intent “file claim” was improperly matched to the natural language query submitted by the first user device. Alternatively, if the first user device accesses the article and uses the website to file the claim, the server may determine that the intent “file claim” was properly matched to the natural language query submitted by the first user device.

[0091] At **906**, the server transmits the natural language query and the intent to a second user device (e.g., the reviewing user device **506**). In some examples, a page (e.g., the page **804**) associated with a given intent (e.g., the intent “process refund”) may include representations of multiple queries (e.g., the queries **808A**, **808B**) that were matched to the intent. The page may be presented in a web browser or in a dedicated application associated with the server.

[0092] At **908**, the server receives, from the second user device, a response indicating whether the natural language query is properly matched to the intent. The response may include a user selection of a first user interface icon (e.g., the “yes” buttons **810A**, **810B**) indicating that the matching (of the query to the intent) is accurate or a user selection of a second user interface icon (e.g., the “no” buttons **812A**, **812B**) indicating that the matching is inaccurate.

[0093] At **910**, the server trains the intent matching engine based on online learning and the response. In some cases, the server stores training data (e.g., the training data **614**), and the response from the second user device is added to the training data. The training data may be used to train the intent matching engine. The training operation may be run in real time when a new response is added to the training data. Alternatively, to reduce the resources used for training, the training operation may be run once every threshold time period (e.g., once per week or once per month) with all the training data received during the time period. In some examples, the training operation may be run whenever a certain number (e.g., 1000, 10,000, 100,000, or 1,000,000) responses are added to the training data. The number may be adjusted based on the number of previously used training data responses. For example, when there are few training data responses, the training operation may be rerun whenever 100 or 1000 new training responses are received. As the intent matching engine is trained on more and more data, and new training data units do not have much impact on its operation, then number may be increased to 100,000 or 1,000,000.

[0094] In some cases, the server processes multiple natural language queries. The server may monitor a number of queries associated with each intent during multiple time periods. The server may generate (and transmit to an agent device or the second user device) a dashboard displaying the number of queries associated with a given intent during all or a subset of the multiple time periods. For example, an insurance company may have the intents “renew policy” and “file claim” and may track the number of times these intents are found in each month. This dashboard can be



used to generate an early indicator that the number of users filing claims is increasing and/or that the number of users renewing their policies is decreasing. Such an early indicator may be useful, for example, in understanding the finances of the insurance company.

[0095] In one example use case, Anna telephones a taxi service and says, “Quiero ir al aeropuerto,” which can be translated to “I want to go to the airport,” in Spanish. A chatbot at a server of the taxi service obtains the query “Quiero ir al aeropuerto,” and matches this query to the intent “request taxi.” The server may store English language intents and may have been trained to match both Spanish language queries and English language queries to the English language intents.

Alternatively, the server may automatically detect the language of the query and automatically translate the query into English prior to matching the query to the intent.

[0096] The server attempts to map a workflow to the intent “request taxi.” Several workflows are stored in a database coupled with the server. The workflows include: an article explaining how to use a mobile application to call a taxi, a text-operated chatbot website interface to call a taxi, an English language audio-operated chatbot interface to call a taxi, and an English and Spanish-speaking human operator who can assist with calling the taxi. Since Anna is accessing the taxi service over the PSTN, the server determines that either the audio-operated chatbot interface or the human operator would be most appropriate. Furthermore, as Anna spoke her query in Spanish, and the human operator speaks Spanish (while the chatbot does not have a Spanish interface), the server determines that the English and Spanish-speaking human operator is an optimal workflow for Anna (from the available workflows). Anna is then connected with the human operator, who assists Anna with ordering the taxi.

[0097] With appropriate notification to Anna, Anna's telephone call, including the query, could be recorded. The recording of the query spoken by Anna, as well as an indication of the intent to which it was matched, could be transmitted to a human agent for review. The human agent may confirm that the query “Quiero ir al aeropuerto,” was correctly matched to the intent “request taxi,” thereby generating training data for the intent matching engine of the server of the taxi service. The training data may be used to train the intent matching engine.

[0098] At a later time, an agent of the taxi service may access a dashboard review the intents that were submitted in each month of the last several months. The dashboard may include the intents “request taxi,” “dispute charge,” and “lost item.” When selecting the “request taxi” intent on the dashboard, the number of queries matching to the intent may be indicated, as well as the source (e.g., telephone or website) and the language of the queries may be indicated. As a result, an agent accessing the dashboard may be able to determine that the number of Spanish language telephone requests for taxis is increasing. This may encourage the agent to recommend hiring additional Spanish-speaking telephone contact center agents and/or creating a Spanish-language voice-operated chatbot for processing taxi requests.

[0099] Some implementations are described below as numbered examples (Example 1, 2, 3, etc.). These examples are provided as examples only and do not limit the other implementations disclosed herein.

[0100] Example 1 is a method, comprising: obtaining a natural language query from a first user device; matching an intent to the natural language query using an intent matching engine, wherein the intent represents predicted data associated with the natural language query; transmitting the natural language query and the intent to a second user device; receiving, from the second user device, a response indicating whether the natural language query is properly matched to the intent; and training the intent matching engine based on a machine learning technique and the response.

[0101] In Example 2, the subject matter of Example 1 includes, transmitting, to the first user device, a workflow output associated with a workflow to which the intent is mapped, wherein the workflow output comprises at least one of: a hyperlink, a webpage, an article, a list of actions, a chatbot-operated interface, or an initiation of a communication session between a human agent assigned to handle the workflow and a user of the first user device; monitoring activity of the first

user device with respect to the workflow output, wherein the intent matching engine is further trained based on the activity of the first user device with respect to the workflow output.

[0102] In Example 3, the subject matter of Examples 1-2 includes, wherein the natural language query is one of a plurality of queries, wherein each query of the plurality of queries is matched to an intent from a set of intents, the method comprising: monitoring a number of queries from the plurality of queries that are associated with each intent during multiple time periods; and transmitting, to the second user device, a dashboard display indicating the number of queries associated with a given intent during a subset of the multiple time periods.

[0103] In Example 4, the subject matter of Examples 1-3 includes, wherein the response comprises a selection of either a first option indicating that the matching is accurate or a second option indicating that the matching is inaccurate.

[0104] In Example 5, the subject matter of Examples 1-4 includes, wherein the intent is selected from a set comprising global intents for multiple entities and custom intents for an entity associated with the second user device.

[0105] In Example 6, the subject matter of Examples 1-5 includes, wherein receiving from the second user device the response indicating whether matching the intent to the natural language query is accurate comprises: transmitting a page associated with the intent to second user device, wherein the page displays multiple queries, including the natural language query that matched to the intent; and receiving the response from the second user device via the page.

[0106] In Example 7, the subject matter of Examples 1-6 includes, determining that the natural language query is in a natural language different from a natural language associated with the intent; and translating the natural language query into the natural language associated with the intent prior to matching the intent to the natural language query.

[0107] In Example 8, the subject matter of Examples 1-7 includes, determining that the natural language query is in a natural language different from a natural language associated with the intent, wherein the intent is matched to the natural language query without translating the natural language query.

[0108] In Example 9, the subject matter of Examples 1-8 includes, wherein the machine learning technique comprises online learning.

[0109] Example 10 is a computer readable medium storing instructions operable to cause one or more processors to perform operations comprising: obtaining a natural language query from a first user device; matching an intent to the natural language query using an intent matching engine, wherein the intent represents predicted data associated with the natural language query; transmitting the natural language query and the intent to a second user device; receiving, from the second user device, a response indicating whether the natural language query is properly matched to the intent; and training the intent matching engine based on a machine learning technique and the response.

[0110] In Example 11, the subject matter of Example 10 includes, the operations comprising: transmitting, to the first user device, an audio or visual output associated with a workflow to which the intent is mapped, wherein the audio or visual output comprises at least one of: an article, a list of actions, a chatbot-operated interface, or an initiation of a communication session between a human agent assigned to handle the workflow and a user of the first user device; monitoring activity of the first user device with respect to the audio or visual output, wherein the intent matching engine is further trained based on the activity of the first user device with respect to the audio or visual output.

[0111] In Example 12, the subject matter of Examples 10-11 includes, wherein the natural language query is one of a plurality of queries, wherein each query of the plurality of queries is matched to an intent from a set of intents, the operations comprising: monitoring a number of queries from the plurality of queries that are associated with each intent during a set of discrete time periods; and transmitting, to the second user device, a dashboard display indicating the number of queries

associated with a given intent during a subset of the discrete time periods.

[0112] In Example 13, the subject matter of Examples 9-12 includes, wherein the response comprises a selection of either a first graphical user interface icon indicating that the matching is accurate or a second graphical user interface icon indicating that the matching is inaccurate.

[0113] In Example 14, the subject matter of Examples 10-13 includes, wherein receiving from the second user device the response indicating whether matching the intent to the natural language query is accurate comprises: transmitting a graphical user interface associated with the intent to second user device, wherein the graphical user interface displays multiple queries, including the natural language query that matched to the intent; and receiving the response from the second user device via the graphical user interface.

[0114] In Example 15, the subject matter of Examples 10-14 includes, the operations comprising: determining that the natural language query is in a first natural language, wherein the intent is associated with a second natural language different from the first natural language; and translating the natural language query into the second natural language prior to matching the intent to the natural language query.

[0115] In Example 16, the subject matter of Examples 10-15 includes, the operations comprising: determining that the natural language query is in a first natural language, wherein the intent is associated with a second natural language different from the first natural language, wherein the intent is matched to the natural language query without translating the natural language query into the second natural language.

[0116] Example 17 is an apparatus, comprising: a memory; and a processor configured to execute instructions stored in the memory to: obtain a natural language query from a first user device; match an intent to the natural language query using an intent matching engine, wherein the intent represents predicted data associated with the natural language query; transmit the natural language query and the intent to a second user device; receive, from the second user device, a response indicating whether the natural language query is properly matched to the intent; and train the intent matching engine based on a machine learning technique and the response.

[0117] In Example 18, the subject matter of Example 17 includes, the processor configured to execute the instructions stored in the memory to: map the intent to a workflow; transmit, to the first user device, an audio or visual output associated with the workflow, wherein the audio or visual output comprises at least one of: an article, a list of actions, a chatbot-operated interface, or an initiation of a communication session between a human agent assigned to handle the workflow and a user of the first user device; monitor activity of the first user device with respect to the transmitted audio or visual output, wherein the intent matching engine is further trained based on the activity of the first user device with respect to the transmitted audio or visual output.

[0118] In Example 19, the subject matter of Examples 17-18 includes, wherein the natural language query is one of a plurality of queries, wherein each query of the plurality of queries is matched to an intent from a set of intents, the operations comprising: monitoring a number of queries from the plurality of queries that are associated with each intent over time; and transmitting, to the second user device, a graphical dashboard display indicating the number of queries associated with a given intent over time.

[0119] In Example 20, the subject matter of Examples 17-19 includes, wherein the response comprises a selection of either a first user interface element indicating that the matching is accurate or a second user interface element indicating that the matching is inaccurate.

[0120] Example 21 is at least one machine-readable medium including instructions that, when executed by processing circuitry, cause the processing circuitry to perform operations to implement of any of Examples 1-20.

[0121] Example 22 is an apparatus comprising means to implement of any of Examples 1-20.

[0122] Example 23 is a system to implement of any of Examples 1-20.

[0123] Example 24 is a method to implement of any of Examples 1-20.

[0124] The implementations of this disclosure can be described in terms of functional block components and various processing operations. Such functional block components can be realized by a number of hardware or software components that perform the specified functions. For example, the disclosed implementations can employ various integrated circuit components (e.g., memory elements, processing elements, logic elements, look-up tables, and the like), which can carry out a variety of functions under the control of one or more microprocessors or other control devices. Similarly, where the elements of the disclosed implementations are implemented using software programming or software elements, the systems and techniques can be implemented with a programming or scripting language, such as C, C++, Java, JavaScript, assembler, or the like, with the various algorithms being implemented with a combination of data structures, objects, processes, routines, or other programming elements.

[0125] Functional aspects can be implemented in algorithms that execute on one or more processors. Furthermore, the implementations of the systems and techniques disclosed herein could employ a number of conventional techniques for electronics configuration, signal processing or control, data processing, and the like. The words “mechanism” and “component” are used broadly and are not limited to mechanical or physical implementations, but can include software routines in conjunction with processors, etc. Likewise, the terms “system” or “tool” as used herein and in the figures, but in any event based on their context, may be understood as corresponding to a functional unit implemented using software, hardware (e.g., an integrated circuit, such as an ASIC), or a combination of software and hardware. In certain contexts, such systems or mechanisms may be understood to be a processor-implemented software system or processor-implemented software mechanism that is part of or callable by an executable program, which may itself be wholly or partly composed of such linked systems or mechanisms.

[0126] Implementations or portions of implementations of the above disclosure can take the form of a computer program product accessible from, for example, a computer-usable or computer-readable medium. A computer-usable or computer-readable medium can be a device that can, for example, tangibly contain, store, communicate, or transport a program or data structure for use by or in connection with a processor. The medium can be, for example, an electronic, magnetic, optical, electromagnetic, or semiconductor device.

[0127] Other suitable mediums are also available. Such computer-usable or computer-readable media can be referred to as non-transitory memory or media, and can include volatile memory or non-volatile memory that can change over time. The quality of memory or media being non-transitory refers to such memory or media storing data for some period of time or otherwise based on device power or a device power cycle. A memory of an apparatus described herein, unless otherwise specified, does not have to be physically contained by the apparatus, but is one that can be accessed remotely by the apparatus, and does not have to be contiguous with other memory that might be physically contained by the apparatus.

[0128] While the disclosure has been described in connection with certain implementations, it is to be understood that the disclosure is not to be limited to the disclosed implementations but, on the contrary, is intended to cover various modifications and equivalent arrangements included within the scope of the appended claims, which scope is to be accorded the broadest interpretation so as to encompass all such modifications and equivalent structures as is permitted under the law.

## Claims

1. A method, comprising: storing a set of intents that are matched to natural language queries based on previous natural language queries and network activities of user devices in response to previous workflow outputs; matching an intent to a natural language query using an intent matching engine, wherein the natural language query is obtained from a first user device within a chatbot, wherein the intent represents predicted data associated with the natural language query, wherein each intent

is mapped to a workflow in a data repository accessed by a server, and wherein the set of intents includes a set of global intents that are globally applicable and a set of custom intents that are only applicable to the server; transmitting, to the first user device within the chatbot, a workflow output associated with the workflow to which the intent is mapped; monitoring network activity within the chatbot of the first user device with respect to the workflow output; and training, using online learning applied to an artificial neural network of the intent matching engine, the intent matching engine based on the monitored network activity of the first user device.

2. The method of claim 1, wherein the workflow output comprises at least one of: a hyperlink, a webpage, an article, or a list of actions.

3. The method of claim 1, wherein the natural language query is one of a plurality of queries, wherein each query of the plurality of queries is matched to an intent from the set of intents, the method comprising: monitoring a number of queries from the plurality of queries that are associated with each intent during multiple time periods; and transmitting, to a second user device, a dashboard display indicating the number of queries associated with a given intent during a subset of the multiple time periods.

4. The method of claim 1, wherein the response comprises a selection of either a first option indicating that the matching is accurate or a second option indicating that the matching is inaccurate.

5. The method of claim 1, wherein the workflow output comprises at least one of: a chatbot-operated interface, or initiating a communication session between a human agent assigned to handle the workflow and a user of the first user device.

6. The method of claim 1, further comprising: transmitting the natural language query and the intent to a second user device in response to matching the intent to the natural language query by transmitting a page associated with the intent to the second user device, wherein the page displays multiple queries, including the natural language query that matched to the intent; and receiving from the second user device, via the page, a response indicating whether matching the intent to the natural language query is accurate.

7. The method of claim 1, comprising: determining that the natural language query is in a natural language different from a natural language associated with the intent; and translating the natural language query into the natural language associated with the intent prior to matching the intent to the natural language query.

8. The method of claim 1, comprising: determining that the natural language query is in a natural language different from a natural language associated with the intent, wherein the intent is matched to the natural language query without translating the natural language query.

9. The method of claim 1, wherein the chatbot is comprised in a chat application accessible by the first user device via a mobile application or a website.

10. One or more non-transitory computer readable media storing instructions operable to cause one or more processors to perform operations comprising: storing a set of intents that are matched to natural language queries based on previous natural language queries and network activities of user devices in response to previous workflow outputs; matching an intent to a natural language query using an intent matching engine, wherein the natural language query is obtained from a first user device within a chatbot, wherein the intent represents predicted data associated with the natural language query, wherein each intent is mapped to a workflow in a data repository accessed by a server, and wherein the set of intents includes a set of global intents that are globally applicable a set of custom intents that are only applicable to the server; transmitting, to the first user device within the chatbot, a workflow output associated with the workflow to which the intent is mapped; monitoring network activity within the chatbot of the first user device with respect to the workflow output; and training, using online learning applied to an artificial neural network of the intent matching engine, the intent matching engine based on the monitored network activity of the first user device.

- 11.** The one or more computer readable media of claim 10, wherein the workflow output comprises at least one of: a hyperlink, a webpage, an article, or a list of actions.
- 12.** The one or more computer readable media of claim 10, wherein the natural language query is one of a plurality of queries, wherein each query of the plurality of queries is matched to an intent from the set of intents, the operations comprising: monitoring a number of queries from the plurality of queries that are associated with each intent during a set of discrete time periods; and transmitting, to a second user device, a dashboard display indicating the number of queries associated with a given intent during a subset of the discrete time periods.
- 13.** The one or more computer readable media of claim 10, wherein the response comprises a selection of either a first graphical user interface icon indicating that the matching is accurate or a second graphical user interface icon indicating that the matching is inaccurate.
- 14.** The one or more computer readable media of claim 10, the operations comprising: transmitting the natural language query and the intent to a second user device in response to matching the intent to the natural language query by transmitting a graphical user interface associated with the intent to the second user device, wherein the graphical user interface displays multiple queries, including the natural language query that matched to the intent; and receiving from the second user device, via the graphical user interface, the response indicating whether matching the intent to the natural language query is accurate.
- 15.** The one or more computer readable media of claim 10, the operations comprising: determining that the natural language query is in a first natural language, wherein the intent is associated with a second natural language different from the first natural language; and translating the natural language query into the second natural language prior to matching the intent to the natural language query.
- 16.** The one or more computer readable media of claim 10, the operations comprising: determining that the natural language query is in a first natural language, wherein the intent is associated with a second natural language different from the first natural language, wherein the intent is matched to the natural language query without translating the natural language query into the second natural language.
- 17.** A system, comprising: one or more memories; and one or more processors configured to execute instructions stored in the one or more memories to: store a set of intents that are matched to natural language queries based on previous natural language queries and network activities of user devices in response to previous workflow outputs; match an intent to a natural language query using an intent matching engine, wherein the natural language query is obtained from a first user device within a chatbot, wherein the intent represents predicted data associated with the natural language query, wherein each intent is mapped to a workflow in a data repository accessed by a server, and wherein the set of intents includes a set of global intents that are globally applicable and a set of custom intents that are only applicable to the server; transmit, to the first user device within the chatbot, a workflow output associated with the workflow to which the intent is mapped; monitor network activity within the chatbot of the first user device with respect to the workflow output; and train, using online learning applied to an artificial neural network of the intent matching engine, the intent matching engine based on the monitored network activity of the first user device.
- 18.** The system of claim 17, wherein the workflow output comprises at least one of: a hyperlink, a webpage, an article, or a list of actions.
- 19.** The system of claim 17, wherein the natural language query is one of a plurality of queries, wherein each query of the plurality of queries is matched to an intent from the set of intents, the one or more processors configured to execute the instructions stored in the one or more memories to: monitor a number of queries from the plurality of queries that are associated with each intent over time; and transmit, to a second user device, a graphical dashboard display indicating the number of queries associated with a given intent over time.
- 20.** The system of claim 17, wherein the response comprises a selection of either a first user

interface element indicating that the matching is accurate or a second user interface element indicating that the matching is inaccurate.

---