

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265614

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Loginov; Nikita et al.

System And Method For Allocating Digital Data According To Value Of Digital Content

Abstract

A novel system for performing valuations of file system objects include a processor configured to execute code and memory configured to store data and the code. The code includes an accumulator configured to accumulate secondary data associated with said file system objects, a data aggregator configured to arrange the secondary data in the memory, and a value generator configured to generate an object value indicative of the value of at least one of the file system objects to a particular entity. In a particular embodiment, the value generator creates a record in a valuation database to associate the object value with the at least one of the file system objects.

Inventors: Loginov; Nikita (Mountain View, CA), Guis; Isabelle Stephanie (Menlo Park, CA), Achalu; Narayanan (Cupertino, CA), Jassal; Amrit (Morgan Hill, CA)

Applicant: Egnyte, Inc. (Mountain View, CA)

Family ID: 1000008578101

Assignee: Egnyte, Inc. (Mountain View, CA)

Appl. No.: 19/029342

Filed: January 17, 2025

Related U.S. Application Data

parent US continuation 18213454 20230623 PENDING child US 19029342

parent US continuation 17243895 20210429 ABANDONED child US 18213454

Publication Classification

Int. Cl.: G06Q30/0201 (20230101); G06F16/11 (20190101); G06F16/14 (20190101); G06F16/17 (20190101); G06F16/28 (20190101); G06Q50/18 (20120101)

U.S. Cl.:

CPC G06Q30/0206 (20130101); G06F16/122 (20190101); G06F16/14 (20190101); G06F16/17 (20190101); G06F16/284 (20190101); G06Q50/184 (20130101);

Background/Summary

RELATED APPLICATIONS [0001] This application is a continuation of co-pending U.S. patent application Ser. No. 18/213,454, filed Jun. 23, 2023 by the same inventors, which is a continuation of U.S. patent application Ser. No. 17/243,895, filed Apr. 29, 2021 by the same inventors, which is a continuation of U.S. patent application Ser. No. 15/937,379, filed Mar. 27, 2018 by the same inventors, which claims the benefit of U.S. Provisional Patent Application No. 62/477,279, filed Mar. 27, 2017 by the same inventors, all of which are incorporated herein by reference in their entireties.

BACKGROUND OF THE INVENTION

Field of the Invention

[0002] This invention relates generally to digital content management, and more particularly to systems and methods of categorizing and determining the value of digital content.

Description of the Background Art

[0003] The storage and management of digital content (e.g., files, folders, etc.) has become commonplace in most organizations today. Knowledge of content value is relevant to multiple IT and business applications, such as defining security policies to protect valuable files, always having a back-up copy of the business's critical information, and identifying and evaluating security “red flags”. Currently, the problem that exists in the content management and analytics space is how to estimate that content value, and do so accurately, in order to drive optimal IT and business decisions.

[0004] There have been attempts to automate the process of estimating file value by means of computer algorithms. Such algorithms scan file content and try to identify important patterns, such as personally identifiable information (credit card numbers, social security numbers) or compliance information (e.g. HIPAA). While this approach works reasonably well for data security applications, it has a few drawbacks when it comes to other types of content value. This is because such an approach cannot estimate the importance of a file to the organization itself in terms of various aspects. At the same time, not all valuable files contain patterns that are identifiable without human involvement, and content scanning does not work for all files. Another problem is that even if content classification works, it cannot always identify fake sensitive content from real, and thus may have a high false positive rate.

[0005] What is needed, therefore, is a system and method for more accurately determining the value of digital content.

SUMMARY

[0006] The inventors have recognized that not all digital content in an organization has the same value and that it is advantageous to quantify the business value of digital content (e.g., files). For example, a particular file (e.g. intellectual property file) may have a much higher security value than a file containing a product catalog. Furthermore, the value(s) of digital content is not necessarily limited to a single category. In other words, the same content can have different values in the context of different paradigms (e.g., security perspective, operational perspective, intellectual property perspective, and so on). For example, a file may have a low business value from a security perspective, but a high business value from an operational perspective due to how frequently it is accessed by members of the organization. As another example, a file containing personally

identifiable information is valuable from the perspective of data security and privacy, but probably not so valuable from an operational perspective. As yet another example, a digital ad for a viral marketing campaign may not be highly valuable from the data security perspective, but may be relatively more valuable from the perspective of instant content access and sharing at scale.

[0007] Various embodiments of the present invention overcome problems associated with the prior art by providing systems and methods of determining values of differing types for file system objects. The systems and methods can be implemented in a local file system or in a distributed cloud file system via a wide-area network.

[0008] An example method for performing valuations of file system objects in a file storage system includes accumulating secondary data associated with the file system objects and arranging the secondary data in a storage device. The secondary data is other than content of the file system objects (i.e., not file system object content). The example method additionally includes generating an object value indicative of the value of at least one of the file system objects to a particular entity, and creating a record in a valuation database. The record associates the object value with the at least one of the file system objects.

[0009] In a particular example method, the step of generating an object value includes generating a plurality of object values. Each object value is indicative of a distinct value of the at least one of the file system objects from a different perspective. For example, one of the plurality of object values can be indicative of a value of the at least one of the file system objects from a data security perspective. As another example, one (or another) of the plurality of object values can be indicative of a value of the at least one of the file system objects from a business continuity perspective. As yet another example, the object value can be expressed in the form of a particular monetary currency (e.g., a dollar value).

[0010] In a particular example method, the secondary data includes data associated with the particular entity but not directly associated with any particular subset of the file system objects. The step of generating an object value includes generating compound data based at least in part on the secondary data accumulated from different sources, and using the compound data to generate the object value. Optionally, the compound data can include user activity data associated with the file system objects. The user activity data is indicative of activity of users of the file storage system with respect to the file system objects (e.g., file access/modification data).

[0011] A more particular example method includes accumulating primary data indicative of content of the individual ones of the file system objects and arranging the primary data in the storage device. The step of accumulating primary data can include scanning the individual ones of the file system objects to obtain the content from the individual ones of the file system objects. The step of accumulating primary data can also include categorizing the individual ones of the file system objects based at least in part on the content of the individual ones of the file system objects.

[0012] The step of accumulating secondary data can include accessing metadata associated with the individual ones of the file system objects. The step of accumulating secondary data can also include accumulating data indicative of an infrastructure of the file storage system. Optionally, the step of accumulating secondary data and/or primary data includes accessing the file system objects of the file storage system over a wide area network.

[0013] An example method additionally includes utilizing at least a portion of the primary data and the secondary data to generate compound data. The compound data can be indicative of relationships between individual ones of the file system objects. In the example method, the step of utilizing at least a portion of the primary data and the secondary data to generate compound data can include estimating the similarity of two or more different ones of the file system objects, based at least in part on the primary data.

[0014] An example method further includes expressing the object value of the at least one of the file system objects relative to object values of other objects of the file system. The step of expressing the object value of the at least one of the file system objects relative to object values of

other objects of the file system can include expressing the object value relative to a base value that is equal to a lowest object value of the object values of the other objects of the file system.

[0015] Example methods can additionally include utilizing the object value(s) of the file system object(s) to inform IT policy changes relating to the particular entity.

[0016] Example valuation systems for performing valuations of file system objects are also disclosed. A particular example valuation system includes a processor configured to execute code and memory configured to store data and the code. The code includes an accumulator, a data aggregator, and a value generator. The accumulator is configured to accumulate secondary data associated with the file system objects. The secondary data is non-object-content data (e.g., does not include file object content). The data aggregator is configured to arrange the secondary data in the memory, and the value generator configured to generate an object value indicative of the value of at least one of the file system objects to a particular entity (e.g., an enterprise). The value generator can also create record(s) in a valuation database to associate the object value(s) with the file system object(s).

[0017] In an example system, the secondary data includes data associated with the particular entity but not directly associated with any particular subset of the file system objects. The value generator can be configured to generate a plurality of different object values for each particular file system object. Each different object value can be indicative of a distinct value associated with the particular file system object from a different perspective. For example, one of the object values of the plurality of object values can be indicative of a value of the file system object(s) from a data security perspective. Another of the plurality of object values can be indicative of a value of the file system object(s) from a business continuity perspective. Optionally, the object value(s) can be expressed in the form of a particular monetary currency (e.g., a dollar value).

[0018] In an example system, the value generator is configured to generate compound data based at least in part on the secondary data accumulated from different sources. The value generator is also configured to utilize the compound data to generate the object value(s).

[0019] The compound data can include user activity data associated with the file system objects. The user activity data (e.g., file access records) is indicative of activity of users of the file storage system with respect to the file system objects.

[0020] In a particular example system, the accumulator is configured to accumulate primary data (e.g., file content) indicative of content of the individual ones of the file system objects. In addition, the data aggregator is configured to arrange the primary data in the storage device. The accumulator is additionally configured to scan the individual ones of the file system objects to obtain the content from the individual ones of the file system objects. The accumulator can be additionally configured to categorize the individual ones of the file system objects based at least in part on the content of the individual ones of the file system objects. The accumulator can also be configured to access metadata associated with the individual ones of the file system objects and/or accumulate data indicative of an infrastructure of the file storage system. Optionally, the example system additionally includes a network adapter configured to access the file system objects of the file storage system over a wide area network (e.g., objects on a distributed cloud file storage system or other file storage system remote from the valuation system).

[0021] The value generator can be configured to utilize at least a portion of the primary data and the secondary data to generate compound data. The compound data can be indicative of relationships between individual ones of the file system objects. For example, the value generator can be additionally configured to estimate the similarity of two or more different ones of the file system objects, based at least in part on the primary data. The value generator can also be configured to express the object value of the at least one of the file system objects relative to object values of others of the file system objects. The value generator can additionally be configured to express the object value relative to a base value that is equal to a lowest object value of the object values of the others of the file system objects.

[0022] An example system additionally includes a valuation enforcer configured to utilize the object value(s) of the file system object(s) to inform IT policy changes relating to the particular entity.

[0023] Another example valuation system, for performing valuations of file system objects, includes an accumulator configured to accumulate secondary data associated with the file system objects. The secondary data is non-object-content data. The example valuation system additionally includes a data aggregator configured to arrange the secondary data in a storage device. The example valuation system additionally includes means for generating an object value indicative of the value of at least one of the file system objects to a particular entity, and for associating the object value with the at least one of the file system objects.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0024] The present invention is described with reference to the following drawings, wherein like reference numbers denote substantially similar elements:

[0025] FIG. 1 is a block diagram of a cloud computing system configured for valuing digital content;

[0026] FIG. 2 is a block diagram showing a remote cloud of FIG. 1 in greater detail;

[0027] FIG. 3 is a block diagram showing valuation services of FIG. 2 in greater detail;

[0028] FIG. 4 is a data flow diagram illustrating the process of determining file values;

[0029] FIG. 5 is a block diagram summarizing an example process flow of the value generating systems of FIG. 1;

[0030] FIG. 6A is a diagram showing a portion of an example database schema of a raw data database of FIG. 4;

[0031] FIG. 6B is a diagram showing another portion of the database schema of FIG. 6A;

[0032] FIG. 6C is a diagram showing yet another portion of the database schema of FIG. 6A;

[0033] FIG. 7A is a diagram showing a portion of an example database schema of a compound data database of FIG. 4;

[0034] FIG. 7B is a diagram showing another portion of the database schema of FIG. 7A;

[0035] FIG. 8 is a diagram showing an example database schema of a file values database of FIG. 4;

[0036] FIG. 9 is a diagram showing an example database schema **900** of a governance policy database of FIG. 4; and

[0037] FIG. 10 is a flow chart summarizing an example method for generating object values corresponding to file system objects in a file storage system.

DETAILED DESCRIPTION

[0038] The present invention overcomes the problems associated with the prior art, by providing a system and method for determining values and value categories for digital content. In the following description, numerous specific details are set forth in order to provide a thorough understanding of the invention (e.g., examples of value categories). Those skilled in the art will recognize, however, that the invention may be practiced apart from these specific details. In other instances, details of well-known computing practices (e.g., database programming, etc.) and components have been omitted, so as not to unnecessarily obscure the present invention.

[0039] FIG. 1 is a block diagram of a cloud computing system **100** configured for valuing digital content. System **100** includes a cloud content storage/computing system (remote cloud) **102**, a local cloud system (local cloud) **104**, and a valuation service software as a service (SaaS) **106**, all intercommunicating via an internetwork **108**. Note that internetwork **108** could be any type of communication network (e.g., the Internet, wide-area network, telecom system, etc.).

[0040] Remote cloud **102** is a distributed remote file storage system and server accessible to clients **110** over internetwork **108**. Remote cloud **102** stores data associated with a particular entity (or a plurality of different entities) (e.g., a business, a cloud customer, etc.) and provides access to the stored data to clients **110**, who are associated with the particular entity. Remote cloud **102** includes a valuation service **112** operative on the data stored therein. Valuation service **112** generates and stores valuation data corresponding to digital objects stored on remote cloud **102**. The valuation data is indicative of the value of individual data objects, from various perspectives, to the particular entity associated with those object on remote cloud **102**.

[0041] Local cloud **104** stores data associated with the particular entity and accessible through a local network **114**. Local clients **116**, having access to local network **114**, can access data stored on local cloud **104**, including data objects, applications, etc. Additional network connected storage device(s) **118** are connected to local network **114**. Network connected storage device(s) **118** provide additional data storage and can be accessed by local clients **116** through local network **114**. Local cloud **104** also includes a valuation service **120**, which generates and stores valuation data corresponding to digital content stored in local cloud **104**, local clients **116**, and/or network connected storage device(s) **118**. In particular embodiments, the data/objects stored on local cloud **104** and remote cloud **102** are synchronized.

[0042] Valuation service SaaS module **106** is a value generation system that is implemented in the form of remote software as a service. SaaS module **108** can be operative on data stored in remote cloud **102**, local cloud **104**, and/or network connected storage devices **118**. SaaS module **106** accesses digital objects (and associated data) stored on remote cloud **102**, local cloud **104**, and/or network connected storage devices **118** through publicly available application programming interfaces (APIs). More information regarding the access of data by SaaS module **106** (as well as remote cloud **102** and local cloud **104**) can be found in U.S. patent application Ser. No. 15/487,947, entitled Hybrid Approach to Data Governance, filed Apr. 14, 2017 by Jassal et al., which is incorporated herein by reference in its entirety.

[0043] Valuation services **112**, **120** and **106** are generally similar in function, but require slight differences in implementation. For example, valuation service **112** has local access to data objects associated with a plurality of different cloud customers and must, therefore, differentiate between data objects belonging to different customers. Valuation service **120** has local access to data objects associated only with the particular entity associated with local network **114**, so has no need to differentiate between data objects associated with different customers, but must be able to access a variety of data sources over local network **114**. Additionally, valuation service SaaS module **106** must access data sources over inter-network **108**. For these reasons, valuation services **112**, **120**, and **106** are similar, but not completely interchangeable. For the sake of brevity, the present invention will be described in more detail with reference to valuation service **112**, and not valuation service **106** or valuation service **120**. However, it will be apparent to those skilled in the art how configure valuation services **120** and **106** in view of the following description and Jassal et al., cited above.

[0044] FIG. 2 is a block diagram showing remote cloud **102** in greater detail. Remote cloud **102** includes a plurality of data governance servers **202(1-S)**, a plurality of multi-tenant storage devices **204**, a governance database **206**, and a wide-area network (WAN) adapter **208**, all interconnected via a local network **210**. Multi-tenant storage devices **204** are, for example, RAID-equipped hard drives, which provide non-volatile storage for data belonging to cloud customers associated with remote cloud **102**. Governance database **206** is a relational database containing data necessary for data governance servers **202** to provide data governance services (including data valuation services) to the cloud customers. The data contained in governance database **206** will be discussed in more detail below. WAN adapter **208** provides Internet access to the other components of remote cloud **102**. Particularly, WAN adapter **208** allows the components of remote cloud **102** to send and receive data objects, software, commands, instructions, etc. from remote systems (e.g. local cloud

104) over the Internet.

[0045] Data governance servers **202** provide data governance services, such as data monitoring, data analytics, etc. for data stored on multi-tenant storage devices **204**. Particularly, data governance servers **202** provide valuation services for digital objects stored on multi-tenant storage devices **204**. Each of data governance servers **202** corresponds to a different cloud customer. Data governance server **20(1)**, for example, corresponds to a first cloud customer (associated with local cloud **104**) and includes a processing unit(s) **212(1)**, working memory **214(1)**, a user interface device(s) **216(1)**, non-volatile memory **218(1)**, a data governance services module **220(1)**, a valuation services module **222(1)**, and a local network adapter **224(1)**, all interconnected and communicating via an internal bus **226(1)**. Processing unit(s) **212(1)** execute code that is transferred into working memory **214(1)** from, for example, non-volatile memory **218(1)** and/or client storage devices **204** to impart functionality to governance server **202(1)**. Working memory **214(1)** stores accessed data, including definitions, domain names, etc., as well as the code that is processed by processing unit(s) **212(1)** to perform the intended functions of data governance server **202(1)**. User interface device(s) **216(1)** is an access interface that allows IT technicians to access data governance server **202(1)** for various reasons, such as software, firmware, and/or hardware updates. Non-volatile memory **218(1)** provides additional non-volatile storage for use by governance server **202(1)**. Data governance services **220(1)** are various software services, running within working memory **214(1)**, for collecting, manipulating, and analyzing data, in order to manage the overall availability, usability, integrity, and security of data belonging to the associated cloud customer. Valuation services **222(1)** include various software services, running within working memory **214(1)**, for collecting, manipulating, and analyzing data, in order to determine the values of data objects (stored in multi-tenant storage devices **204**) to the associated cloud customer. Local network adapter **224(1)** provides a network connection between data governance server **202(1)** and local network **210** and, therefore, WAN adapter **208**, which provides a connection to the Internet **108**. Although only data governance server **202(1)** is shown in detail, it should be understood that data governance server **202(1)** is substantially similar to data governance servers **202(2-S)**, except that any of data governance servers **202** can correspond to different cloud clients and, therefore, can be configured differently to utilize different data, applications, network connections, etc. Optionally, although data governance servers **202(1-S)** are shown as separate machines, data governance servers **202(1-S)** can be implemented with virtual machines running on one or more servers.

[0046] The valuation services implemented on local cloud **104** or valuation service SaaS **106** can utilize hardware differing from data governance servers **202(1-S)**. For example, the valuation services located at local cloud **104** can be implemented on a personal computer, a network connected storage device, or a network attached server, or can be distributed across a plurality of similar devices. Additionally, some software aspects may differ. For example, the valuation services implemented via valuation service SaaS **106** require functionality for accessing data stored on remote storage devices through the Internet. Such variations will be apparent to those skilled in the art, particularly in view of this disclosure.

[0047] FIG. **3** is a block diagram showing valuation services **222(1)**. Valuation services **222(1)** can be software modules running in working memory **214(1)** and include a data aggregator **302**, a compound data generator **304**, a file value generator **306**, and a policy/value reviser **308**. Data aggregator **302** receives raw data relevant to the valuation process via local network adapter **224(1)** and, optionally, WAN adapter **208** and organizes the data into a relational database structure stored in governance database **206**. Compound data generator **304** accesses the data stored in governance database **206** to generate compound data and organize the compound data into a relational database structure stored in governance database **206**. The compound data is indicative of relationships between various data objects, users, user groups, etc. that are useful for determining the values of data objects. File value generator **306** utilizes data generated by data aggregator **302** and compound

data generator **304** to determine the values of data objects and organize the values in a relational database structure stored in governance database **206**. Policy/value reviser **308** utilizes machine learning and user feedback to review and revise object values and related data policy and to improve the valuation process.

[0048] Although shown as software modules in working memory **214(1)** being executed by a processor, the functionality of valuation services **222** can be implemented with software, hardware, firmware, or any combination thereof. In addition, any of the methods disclosed herein can be implemented with an electronically-readable medium having code embodied therein for causing an electronic device to carry out the relevant function(s)/method(s).

[0049] FIG. **4** is a data flow diagram illustrating the process of determining file values. Data aggregator **302** retrieves data from data source(s) **402**, which can include multi-tenant storage devices **204** and/or data stored on local cloud **104**, available via the Internet, etc. The received data includes, but is not limited to, file content, file metadata, file access data, user/user-group data, organizational data (related to the cloud customer), and/or infrastructure data. Data aggregator **302** organizes the data into a raw data database **404**. Database **404** is a relational database that is query-able and accessible to compound data generator **304**, file value generator **306**, and policy/value reviser **308**. The data structures contained in database **404** are discussed below, in detail, with reference to FIGS. **6A-6C**. Compound data generator **304** queries the raw data in database **404** in order to generate compound data based on the raw data. The compound data includes, but is not limited to, file similarity data, statistical data, user quality data, and/or relative file/user metrics. Compound data generator **304** organizes the data into a compound data database **406**. Database **406** is a relational database that is query-able and accessible to file value generator **306** and policy/value reviser **308**. The data structures contained in database **406** are discussed below, in detail, with reference to FIGS. **7A-7B**. Databases **404** and **406** contain all the necessary data to generate values of file system objects stored on multi-tenant storage devices **204** and/or local cloud **104**.

[0050] File value generator **306** generates file values for at least a portion of the file objects owned by a cloud customer associated with remote cloud **102**. File values are, for example, numerical representations of the values of files to the cloud customer (e.g. a business) from various perspectives. Each valued file has at least one associated file value, but can have many. For example, each file object can have a value from a data security perspective as well as a business continuity perspective (i.e. how important the file is for the everyday operations of the business). File value generator **306** accesses data stored in databases **404** and **406**, in order to gather data associated with each file object. The gathered data is indicative of the value of each file object from the various perspectives (i.e. value categories). For example, a file that is accessed often by a large number of users in the business (e.g. a pricing table) will be valued highly in the business continuity value category, while a file that is accessed by only a limited number of users in a select group will be valued highly in the data security value category. File value generator **306** generates the values for each file from a particular perspective by weighting each piece of data based on relevance to the particular perspective. For example, from a business continuity perspective, how often an object is accessed and by how many users is of primary importance, whereas the content, location, and/or size of the file are of lesser or no importance. File value generator **306** then organizes all the file values in a file value database **408**. Database **408** is a relational database that is query-able and accessible by policy/value reviser **308**. In addition to the file valuations, database **408** includes data useful to provide file values in terms of monetary currency. The data structures contained in database **408** are discussed below, in detail, with reference to FIG. **8**.

[0051] Once the file values have been generated and stored, they are accessed by policy/value reviser **308** to inform information technology (IT) policies and practices of the business, as well as to evaluate and refine the valuations themselves. Policy/value reviser **308** utilizes user feedback, which it gathers through user input and stores in database **408**. Policy/value reviser **308** utilizes the

user feedback, as well as data from databases **404** and **406** to determine whether file values should be raised, lowered, or maintained at their current values. Additionally, policy/value reviser **308** utilizes the file values to influence decisions regarding data governance. Particularly, policy/value reviser **308** revises policies stored in a governance policy database **410**.

[0052] For example, policy/value reviser **308** can alter a policy that causes particular files to be stored with more or fewer backups, in more or less expensive storage devices, etc. Database **410** is a relational database that is query-able and can be accessible to other components of remote cloud **102** and/or local cloud **104**. Databases **404**, **406**, **408**, and **410** represent various portions of governance database **206** (FIG. 2).

[0053] In addition to user input, data aggregator **302**, compound data generator **304**, file value generator **306**, and policy/value reviser **308** utilize machine learning to refine the process of generating file values. For example, compound data generator **304** can track how often various forms of compound data are utilized in generating file values and generate more or less similar data in the future. As another example, policy/value reviser **308** can recognize how file-level attributes change with changes to file values, and utilize this information in generating future file values. These examples are not intended to be limiting, as the components of valuation services **222(1)** can utilize any relevant machine learning technique, including those yet to be invented.

[0054] FIG. 5 is a block diagram summarizing an example process flow **500** of value generating systems **112**, **120**, and/or **106**. In a first part **502** of process flow **500**, raw data is gathered from several different sources, including a third party **504**, an associated organization **506** and its users, and content infrastructure **508**. Third party **504** can be, by way of non-limiting example, a trade association, a regulatory body, a news outlet, etc. Data gathered from third party **504** (and other third parties) includes, at least, data about organization **506** and/or the associated industry, such as industry type, average company size within the industry, average revenue within the industry, etc. Organization **506** is the cloud customer utilizing the valuation services. Data can be gathered from storage devices, communications, publications, etc., which are owned by or associated with organization **506**. The gathered data can include, for example, data about the organization, such as employment statistics, costs, revenue, etc. The users of organization **506** are the people within the organization (e.g. employees) that utilize user accounts to view, create, copy, delete, or otherwise access or modify data objects belonging to organization **506**. Data gathered from the users includes user data, such as user permissions, user groups, etc. User activity data, such as object access statistics, is also gathered whenever a user of organization **506** interacts with data stored on content infrastructure **508**. Content infrastructure **508** includes, for example, the data storage infrastructure utilized to store and access data owned by organization **506**. Data gathered from content infrastructure **508** includes data about the content infrastructure itself, such as device types and locations, as well as data about the data objects stored within content infrastructure **508**, such as file-and/or folder-level metadata and/or content data (i.e. content of the files/objects themselves). All of the data gathered in part **502** of process flow **500** can be relevant to the valuations of individual files or groups of files, in either absolute terms, relative terms, or as expressed in the form of a monetary currency. For example, the revenue of organization **506** is important, but not necessary, for expressing the value of files in the form of a monetary currency.

[0055] In a second part **510** of process flow **500**, the raw data gathered in first part **502** is aggregated and organized into a relational database structure that can be accessed and modified for the purposes of generating compound data and file values. In a third part **512** of process flow **500**, the aggregated raw data is utilized to generate compound data, which is indicative of, among other things, relationships between data objects, users, user groups, entity type, and so on. Third part **512** of process flow **500** includes, for example, a file similarity estimation **514**, a user/group-level data aggregation **516**, and a folder-level data aggregation **518**. File similarity estimation **514** utilizes file content data to determine the similarity between different files in order to, for example, allow IT personnel to locate different instances of the same file, protect sensitive content, etc. User/group

level data aggregation **516** utilizes user data as well as user activity data to generate statistical data relating users, groups of users, and data objects, in order to aid IT personnel in recognizing and reacting to suspicious activity, data access trends, etc. Folder level aggregation **518** utilizes file and folder metadata to generate folder-level data indicative of folder contents within the file system hierarchy of organization **506**. Folder-level data is utilized to generate aggregate values of folders containing pluralities of files, the aggregate folder values being important from a data storage, access, and security standpoint. Folder-level data is useful in protecting valuable data without compromising the overall structure of the file system hierarchy. The data generated in part **512** of process flow **500** is useful in generating accurate file valuations, by providing complex file system information in an easily digestible format.

[0056] In a fourth part **520** of the process flow **500**, the compound data is utilized to generate file values in a plurality of value categories **522**. Value categories **522** include, by way of non-limiting example, a business continuity category **524**, a profitability category **526**, a security/IP protection category **528**, a regulation/sensitive content category **530**, a sharing category **532**, and a reputation category **534**. A file value in business continuity category **524** is indicative of the impact of a corresponding file on the organization's day-to-day operations. For example, a file that is accessed often by a large number of users will likely have a high value in business continuity category **524**. A file value in profitability category **526** is indicative of the impact of a corresponding file on cost and revenue. For example, a file that contains information that is important for minimizing costs or maximizing revenue will likely have a high value in profitability category **526**. A file value in security/IP protection category **528** is indicative of the value of a corresponding file from a security/IP protection perspective. For example, a file that includes information about an undisclosed invention will likely have a high value in security/IP protection category **528**. A file value in regulation/sensitive content index **530** is indicative of the value of a corresponding file from a regulation/sensitive content perspective. For example, a file that includes information protected by The Health Insurance Portability and Accountability Act (HIPAA) will likely have a high value in regulation/sensitive content category **530**. A file value in sharing category **532** is indicative of the value of a corresponding file from a sharing perspective. For example, a file that includes an ad intended for use in a viral marketing campaign will likely have a high value in sharing category **532**. A file value in reputation category **534** is indicative of the value of a corresponding file from a reputation perspective. For example, a file containing information that could damage the reputation of organization **506** will likely have a high value in reputation category **534** (indicating that the file should be protected). Each file can have a value, and in fact different values, in any number of value categories **522**. For example, many files might have a high value in both business continuity category **524** and profitability category **526**, as a file that is important for day-to-day operations will likely also be important for maximizing profits. As another example, a file that contains confidential patient information may have a high value in both security/IP protection category **528** and regulation/sensitive content category **530**, but have a low value (or no value at all) in sharing category **532**.

[0057] In a fifth part **536** of process flow **500**, the file values are utilized to inform organizational policies related to data storage, access, availability, etc. The file values are utilized to define efficient IT, infrastructure, and operational workflow policies, set up alerts and notifications, and create marketing campaigns. For example, files having high values in security/IP protection category **528** can be stored in more secure forms of storage (e.g., behind firewalls, encrypted, and so on), whereas files having high values in sharing category **532** can be stored in more accessible forms of storage. Additionally, files having high values in business continuity category **524** can be stored in multiple locations close to the users accessing them. As another example, files with high values in profitability category **526** can have more backups than other files. As yet another example, files having high values in security/IP protection category **528** and/or regulation/sensitive content category **530** can be protected with alerts or notifications to IT personnel upon any access.

As yet another example, files having high values in sharing category **532** can be utilized in creating marketing campaigns for widespread sharing and access. FIGS. **6A-6C** are diagrams illustrating an example database schema **600** of raw data database **404**. FIG. **6A** shows a portion of database schema **600**, including an industry/third party data table **602**, an organizational data table **604**, a user data table **606**, a group data table **608**, and a group/user data table **610**.

[0058] Each record in industry/third party data table **602** includes an industry_entry_id field **612**, an industry_id field **614**, a regulatory_data field **616**, a comp_size_avg field **618**, and an other field **620**. A record is created in industry/third party data table **602** each time a new piece of data corresponding to a particular industry (or more than one industry) is received by data aggregator **302**.

[0059] Industry_entry_id field **614** is the key field of industry/third party table **602** and an alphanumeric string, which uniquely identifies each record of table **602**. Industry_id field **614** contains data that uniquely identifies the industry (e.g., healthcare, cloud computing, law, etc.), or a group of industries, corresponding to the record. Regulatory_data field **616** contains data indicative of a piece of regulatory data, such as a new regulatory requirement (e.g., certain documents are client privileged). Comp_size_avg field **618** contains data indicative of the average size of companies in the industry specified by industry_id field **614**. Typically a record having an entry in comp_size_avg field **618** is created initially and updated on a regular basis (e.g., once a year). In other entries comp_size_avg field **618** can have a null value. Other field **620** represents one or more additional fields that can contain data indicative of any other information that might be relevant to a particular industry.

[0060] Each record in organizational data table **604** includes a client_id field **622**, an industry_id field **624**, a company_revenue field **626**, a company_size field **628**, and an other field **630**. A record is created in organizational data table **604** for each cloud customer utilizing the valuation services.

[0061] Client_id field **622** is the key field of organizational data table **604** and an alphanumeric string, which uniquely identifies each cloud customer utilizing the valuation services. Industry_id field **624** contains data indicative of the industry that the cloud customer belongs to.

[0062] Industry_id field **624** is analogous to industry_id field **614** of table **602**; identical alphanumeric strings in industry_id field **624** and industry_id field **614** refer to the same industry. Industry_id field **624** and industry_id field **614** create a many-to-many relationship between organizational data table **604** and industry/third party data table **602**, because each industry can have many entries in table **604**, as well as many entries in table **602**. Company_revenue field **626** contains data indicative of the cloud customer's revenue. Company_size field **628** contains data indicative of the size of the cloud customer (e.g. number of employees). Other field **630** contains data indicative of any other data that might be relevant to a particular cloud customer.

[0063] Each record in user data table **606** includes a user_id field **632**, a client_id field **634**, a user_permissions field **636**, and an other field **638**. A record is created in user data table **606** for each user (i.e. a user profile with access to the data objects being valued) in a cloud customer's file system.

[0064] User_id field **632** is the key field of user data table **606** and an alphanumeric string, which uniquely identifies each user. Client_id field **634** contains data indicative of the cloud customer associated with the user. Client_id field **634** is analogous to client_id field **622** of table **604**. Client_id field **634** and client id field **622** create a many-to-one relationship between table **606** and table **604**, because each cloud customer is associated with many users, but each user is only associated with one cloud customer. User_permissions field **636** contains data indicative of object permissions associated with the user. (It should be noted that user permissions can, alternatively, be entered in a completely separate table, which relates users, objects, and permissions on an individual basis.) Other field **638** represents one or more additional fields that can contain data indicative of any other data that might be relevant to a particular user.

[0065] Each record in group data table **608** includes a group_id field **640**, a client_id field **642**, a

group_permissions field **644**, and an other field **646**. A record is created in group data table **608** each time a group is created in a cloud customer's file system.

[0066] Group_id field **640** is the key field of group data table **608** and an alphanumeric string, which uniquely identifies each group. Client_id field **642** contains data indicative of the cloud customer associated with the group. Although not shown, client_id field **642** creates a many-to-one relationship between group data table **608** and organizational data table **604**, because each cloud customer is associated with many groups, but each group is associated with only one cloud customer. Group_permissions field **644** contains data indicative of object permissions associated with the group. (It should be noted that group permissions can also be entered in a separate table, which relates groups, objects, and permissions on an individual basis.) Other field **646** represents one or more additional fields that can contain data indicative of any other data that might be relevant to a particular group.

[0067] Each record in group/user data table **610** includes a group_entry_id field **648**, a user_id field **650**, a group_id field **652**, and an other field **654**. A record is created in group/user data table **610** each time a user is added to a group.

[0068] Group_entry_id field **648** is the key field of group/user data table **610** and an alphanumeric string, which uniquely identifies each entry in table **610**. User_id field **650** and group_id field **652** together indicate that a specified user belongs to a specified group. User_id field **650** is analogous to user_id field **632** of table **606**. Therefore, user_id field **650** and user_id field **632** create a many-to-one relationship between group/user data table **610** and user data table **606**, because each user is associated with many group/user entries, but each group/user entry is only associated with a single user. Group_id field **652** is analogous to group_id field **640** of table **608**. Therefore, group_id field **652** and group_id field **640** create a many-to-one relationship between group/user data table **610** and group data table **608**, because each group is associated with many group/user entries, but each group/user entry is only associated with a single group. Other field **654** represents one or more additional fields that can contain data indicative of any other data that might be relevant to a particular group/user data entry.

[0069] FIG. **6B** shows another portion of database schema **600**, including user data table **606**, an object metadata table **656**, an object access data table **658**, and a device data table **660**. User data table **606** is shown here (as well as in FIG. **6A**) in order to clearly illustrate the relationships between the tables.

[0070] Each record in object metadata table **656** includes an object_id field **662**, a client_id field **664**, a device_id field **666**, an extension field **668**, a size field **670**, a creation_time field **672**, a last_mod_time field **674**, a security_policy field **676**, and an other field **678**. A record is created in object metadata table **656** each time an object is created in the cloud customer's file system. Object_id field **662** is the key field of object metadata table **656** and contains data that uniquely identifies each file object. Client_id field **664** contains data that identifies the cloud customer associated with object. Client_id field **664** creates a relationship between object metadata table **656** and any other table containing a client_id field. Device_id field **666** contains data that identifies a device on which the object is stored. Extension field **668** contains data that identifies a file extension of the object. Size field **670** contains data that identifies a size of the object. Folders have a default entry of zero (0) in size field **670**. Creation_time field **672** contains data that identifies a date and time that the file was created. Last_mod_time field **674** contains data that identifies a date and time that the file was last modified. Security_policy field **676** contains data that indicates a particular security policy (or multiple policies) that is applied to the object. Other field **678** represents one or more additional fields that can contain data indicative of any other data that might be relevant to a particular data object.

[0071] Each record in object access data table **658** includes an access_event_id field **680**, a user_id field **681**, an object_id field **682**, an access_type field **683**, an access_time field **684**, and an other field **685**. A record is created in object access data table **658** each time a data object is accessed on

the cloud customer's file system.

[0072] Access_event_id field **680** is the key field of object access data table **658** and an alphanumeric string, which uniquely identifies each object access record. User_id field **681** is analogous to user_id field **632** of table **606**. Therefore, user_id field **681** and user_id field **632** create a many-to-one relationship between object access data table **658** and user data table **606**, because each user is associated with many access events, but each access event is only associated with a single user. Object_id field **682** is analogous to object_id field **662** of table **656**. Therefore, object_id field **682** and object_id field **662** create a many-to-one relationship between object access data table **658** and object metadata table **656**, because each object is associated with many access events, but each access event is only associated with a single object.

[0073] Access_type field **683** contains data that identifies how the object was accessed (e.g., READ, WRITE, DELETE, etc.). Access_time field **684** contains data that identifies a date and time that the access occurred. Other field **685** represents one or more additional fields that can contain data indicative of any other data that might be relevant to a particular object access event.

[0074] Each record in device data table **660** includes a device_id field **686**, a client_id field **687**, a device_type field **688**, a device_location field **689**, a security_policy field **690**, a backup_number field **691**, and an other field **692**. A record is created in device data table **660** for each storage device in the cloud customer's file system.

[0075] Device_id field **686** is the key field of device data table **660** and an alphanumeric string, which uniquely identifies each storage device. Device_id field **686** is analogous to device_id field **666** of table **656**. Therefore, device_id field **686** and device_id field **666** create a one-to-many relationship between device data table **660** and object metadata table **656**, because each storage device is associated with many objects, but each object is only associated with a single storage device. Client_id field **687** contains data that identifies the cloud customer associated with the storage device (and is analogous to other client_id fields in other data tables). Device_type field **688** contains data that identifies a type (e.g., SSD, HD, Flash, Disk, etc.) of the storage device. Device_location field **689** contains data that identifies a location (e.g., Chicago office, accounting department, a domain name, etc.) of the storage device. Security_policy field **690** contains data indicative of a security policy that is applied to the storage device. Backup_number field **691** contains data indicative of the number of backups of the storage device that exist. Other field **692** represents one or more additional fields that can contain data indicative of any other data that might be relevant to a particular storage device.

[0076] FIG. 6C shows yet another portion of database schema **600**, including organizational data table **604**, object metadata table **656**, and an object content data table **693**. Organizational data table **604** and object metadata table **656** are shown here (as well as in FIGS. 6A and 6B, respectively) in order to clearly illustrate the relationships between the tables.

[0077] Object metadata table **656** and organizational data table **604** are related by client_id field **664** and client_id field **622**. Client_id field **664** and client_id field **622** create a many-to-one relationship between object metadata table **656** and organizational data table **604**, because each cloud customer is associated with many objects, but each object is only associated with a single cloud customer.

[0078] Object content data table **693** contains a content_record_id field **694**, an object_id field **695**, a content field **696**, a content_type field **697**, and an other field **698**. A record is created in object content data table **693** each time data aggregator **302** receives data indicative of the content of a file system object.

[0079] Content_record_id field **694** is the key field of object content data table **693** and an alphanumeric string, which uniquely identifies each record in table **693**. Object_id field **695** contains data indicative of the object associated with the record. Object_id field **695** is analogous to object_id field **662** of table **656**. Therefore, object_id field **695** and object_id field **662** create a many-to-one relationship between object content data table **693** and object metadata table **656**,

because each object is associated with many content records, but each content record is only associated with a single object. Content field **696** contains data indicative of at least a portion of the content of the object. Content_type field **697** contains data indicative of a type (e.g., non-disclosure agreement (NDA), HIPAA, financial, social security number, etc.) of the content identified by content field **696**. Other field **698** represents one or more additional fields that can contain data indicative of any other information that might be relevant to a particular object content data entry. [0080] FIGS. 7A and 7B are diagrams illustrating an example database schema **700** of compound data database **406**. FIG. 7A shows a portion of database schema **700**, including object metadata table **656**, an object similarity data table **702**, a compound object access data table **704**, and a relative object metrics table **706**. It should be noted that although object metadata table **656** is shown in both raw data database **404** and compound data database **406**, it is only saved once in governance database **206** (aside from any backup copies).

[0081] Object similarity data table **702** includes a similarity_record_id field **708**, an object_id field **710**, a similar_object_id field **712**, a similarity_rating field **714**, and an other field **716**. A record is created in object similarity data table **702** each time compound data generator **304** performs a comparison of two data objects or, alternatively, if compound data generator **304** finds that two objects have a similarity higher than a threshold similarity. Compound data generator **304** utilizes data from object metadata table **656** and object content data table **693** to generate object similarity data table **702**.

[0082] Similarity_record_id **708** is the key field of object similarity data table **702** and an alphanumeric string, which uniquely identifies each record in object similarity data table **702**. Object_id **710** contains data uniquely identifying a particular data object. Object_id field **710** is analogous to object_id field **662** and, therefore, creates a many-to-one relationship between object similarity data table **702** and object metadata table **656**, because each object similarity record corresponds to only one data object, but each data object can correspond to many object similarity records. Similar_object_id field **712** contains data uniquely identifying a particular data object that has been compared to the data object identified by object_id field **710** to determine similarity. For each object comparison there will be two records in object similarity data table **702**: one from the perspective of the object identified by object_id field **710** and another from the perspective of the object identified by similar_object_id field **712**. Similarity_rating field **714** contains data indicative of the similarity (e.g., 50%, 100%, etc.) between the data objects. Other field **716** represents one or more additional fields that can contain data indicative of any other data that might be relevant to an object similarity record.

[0083] Compound object access data table **704** includes an object_id field **718**, an access_frequency field **720**, a perm_group_high field **722**, a perm_group_low field **724**, a perm_user_high field **726**, a perm_user_low field **728**, an object_availability field **730**, and an other field **732**. A record is created in compound object access data table **704** for each data object in the cloud customer's file system. Compound data generator **304** utilizes data from user data table **606**, group data table **608**, object metadata table **656** and object access data table **658** to generate compound object access data table **704**.

[0084] Object_id field **718** is the key field of compound object access data table **704** and an alphanumeric string, which uniquely identifies each data object. Object_id field **718** and object_id field **662** of object metadata table **656** create a one-to-one relationship between records in compound object access data table **704** and object metadata table **656**, because each compound object access record corresponds to only one object and each object corresponds to only one compound object access record. Access_frequency field **720** contains data indicative of the frequency (e.g. **100** times per day) that the data object identified by object_id field **718** is accessed. Perm_group_high field **722** contains data indicative of the group permitted to access the object that has the highest security clearance (i.e. the most relaxed permissions). Perm_group_low field **724** contains data indicative of the group permitted to access the object that has the lowest security

clearance (i.e. the strictest permissions). Perm_group_low field **724** is particularly important for determining data security values, because it is indicative of how secure the object is. For example, if perm_group_low field **724** indicates that executives or IT personnel are the lowest group permitted to access, the corresponding object is most likely highly confidential. Perm_user_high field **726** contains data similar to perm_group_high field **722**, but identifies the user, rather than the group, permitted to access having the highest security clearance. Similarly, perm_user_low field **728** contains data similar to perm_group_low field **724**, but identifies the user permitted to access having the lowest security clearance. Object_availability field **730** contains data indicative of the accessibility of the data object. For example, object_availability field **730** can indicate how many users in the organization have access to the object, how long it takes for the average user to access the object, how difficult the object is to find, etc. Other field **732** represents one or more additional fields that can contain data indicative of any other information that might be relevant to a compound object access record.

[0085] Relative object metrics table **706** includes an object_id field **734**, a rel_size field **736**, rel_user_perms field **738**, a rel_group_perms field **740**, a rel_access_freq field **742**, and an other field **744**. A record is made in relative object metrics table **706** for each object in object metadata table **656**. Compound data generator **304** utilizes user data table **606**, group data table **608**, object metadata table **656**, and compound object access data table **704** to generate relative object metrics table **706**.

[0086] Object_id field **734** is the key field of relative object metrics table **706** and an alphanumeric string, which uniquely identifies each data object. Object_id field **734** and object_id field **662** of object metadata table **656** create a one-to-one relationship between records in relative object metrics table **706** and object metadata table **656**, because each relative object metrics record corresponds to only one object and each object corresponds to only one relative object metrics record. Rel_size field **736** contains data that is indicative of the size of the object identified by object_id field **734** relative to the rest of the objects in the cloud customer's file system. For example, rel_size field **736** can indicate that the object is in the 50th percentile for object size. Rel_user_perms field **738** contains data indicative of the user permissions associated with the object relative to the rest of the objects. For example, rel_user_perms field **738** can indicate that the object is in the 50th percentile for number of open user permissions. Rel_group_perms field **740** contains data indicative of the group permissions associated with the object relative to the rest of the objects. For example, rel_group_perms field **740** can indicate that the object is in the 50th percentile for number of open group permissions. Rel_access_freq field **742** contains data indicative of the frequency the data object is accessed relative to the rest of the objects. For example, rel_access_freq field **742** can indicate that the object is in the 50th percentile for access frequency. Other field **744** represents one or more additional fields that can contain data indicative of any other information that might be relevant to the relative metrics associated with a file system object.

[0087] FIG. 7B shows a portion of database schema **700**, including user data table **606**, a user security data table **746**, a compound user access data table **748**, and a relative user metrics table **750**. It should be noted that although user data table **606** is shown in both raw data database **404** and compound data database **406**, it is only saved once in governance database **206** (aside from any backup copies).

[0088] User security data table **746** includes a user_id field **752**, a security_clearance field **754** and an other field **756**. A record is created in user security data table **746** for each user of the cloud customer's file system. Compound data generator **304** utilizes user data table **606**, group data table **608**, and user/group data table **610** in order to generate user security data table **746**.

[0089] User_id field **752** is the key field of user security data table **746** and an alphanumeric string, which uniquely identifies a particular user. User_id field **752** and user_id field **632** of user data table **606** create a one-to-one relationship between records in user security data table **746** and user

data table **606**, because each user security record corresponds to only one user and each user corresponds to only one user security record. Security_clearance field **754** contains data indicative of a security clearance of the user identified by user_id field **752**. For example, a user with a high level of access might have a “top-secret” clearance identified by security_clearance field **754**. Other field **756** represents one or more additional fields that can contain data indicative of any other information that might be relevant to the security data associated with a file system user. [0090] Compound user access data table **748** includes a user_access_rec_id field **758**, a user_id field **760**, a object_id field **762**, a user_access_freq field **764**, a user_access_reg field **766**, and an other field **768**. A record is created in compound user access data table **748** for each user-object pair in the cloud customer's file system. Compound data generator **304** utilizes user data table **606**, object metadata table **656**, and object access data table **658** to generate compound user access data table **748**.

[0091] User_access_rec_id field **758** is the key field of compound user access data table **748** and an alphanumeric string, which uniquely identifies each record in compound user access data table **748**. User_id field **760** contains data indicative of a particular user. User_id field **760** and user_id field **632** of user data table **606** create a many-to-one relationship between records in compound user access data table **748** and user data table **606**, because each compound user access record corresponds to only one user, but each user corresponds to many (one for each file system object) compound user access records. Object_id field **762** contains data indicative of a particular data object. Object_id field **762** is analogous to object_id field **662** of object metadata table **656**. User_access_freq field **764** contains data indicative of how often (e.g. **10** times per day, week, year, etc.) the user identified by user_id field **760** accesses the object identified by object_id field **762**. User_access_reg field **766** contains data indicative of the regularity (e.g., daily, weekly, etc.) at which the user accesses the object. Other field **768** represents one or more additional fields that can contain data indicative of any other information that might be relevant to the access data associated with a particular user-object pair.

[0092] Relative user metrics table **750** includes a rel_user_record_id field **770**, a user_id field **772**, an object_id field **774**, a rel_user_access_freq field **776**, a rel_user_access_reg field **778**, and an other field **780**. A record is created in relative user metrics table **750** for each user-object pair in the cloud customer's file system. Compound data generator **304** utilizes user data table **606**, object metadata table **656**, and compound user access data table **748** to generate relative user metrics table **750**.

[0093] Rel_user_record_id **770** is the key field of relative user metrics table **750** and an alphanumeric string, which uniquely identifies each record in relative user metrics table **750**. User_id field **772** contains data indicative of a particular user. User_id field **772** and user_id field **632** of user data table **606** create a many-to-one relationship between records in relative user metrics table **750** and user data table **606**, because each relative user metrics record corresponds to only one user, but each user corresponds to many (one for each file system object) relative user metrics records. Object_id field **774** contains data indicative of a particular data object. Object_id field **774** is analogous to object_id field **662** of object metadata table **656**. Rel_user_access_freq field **776** contains data indicative of how often the user identified by user_id field **772** accesses the object identified by object_id field **774** relative to the rest of the users. For example, rel_user_access_freq **776** might indicate that the user accesses the object more often than 50% of the other users. Rel_user_access_reg field **778** contains data indicative of the regularity at which the user accesses the object relative to the rest of the users. For example, rel_user_access_reg field **778** might indicate that the user accesses the object more regularly than 50% of the other users. Other field **780** represents one or more additional fields that can contain data indicative of any other information that might be relevant to the relative user metrics associated with a particular user-object pair.

[0094] FIG. **8** shows a database schema **800**, including object metadata table **656**, an object values

table **802**, a valuation feedback data table **804**, and a currency value data table **806**. It should be noted that, although object metadata table **656** is shown in all three of raw data database **404**, compound data database **406**, and file value database **408**, it is only saved once in governance database **206** (aside from any backup copies).

[0095] Object values table **802** includes an object_id field **808**, a continuity_value field **810**, a profitability_value field **812**, an IP_protection_value field **814**, a security_value field **816**, a regulation_value field **818**, a sharing_value field **820**, a reputation_value field **822**, and an other field **824**. A record is created in object values table **802** for each file system object that is valued. File value generator **306** utilizes data in database schema **600** and database schema **700** as well as valuation feedback data table **804** and currency value data table **806** in order to generate object values table **802**.

[0096] Object_id field **808** is the key field of object values **802** and an alphanumeric string, which uniquely identifies a particular file system object. Object_id field **808** contains data indicative of a particular data object. Object_id field **808** and object_id field **662** of object metadata table **656** create a one-to-one relationship between records in object values table **802** and object metadata table **656**, because each object values record corresponds to only one object and each object corresponds to only one object values record. Continuity_value field **810** contains data indicative of a value of the object identified by object_id field **808** from a business continuity perspective. Profitability_value field **812** contains data indicative of a value of the object from a profitability perspective. IP_protection_value field **814** contains data indicative of a value of the object from a IP protection/sensitivity perspective. Security_value field **816** contains data indicative of a value of the object from a data security perspective. Regulation_value field **818** contains data indicative of a value of the object from a regulation perspective. Sharing_value field **820** contains data indicative of a value of the object from a sharing perspective. Reputation_value field **822** contains data indicative of a value of the object from a reputation perspective. Other field **824** represents one or more additional fields that can contain data indicative of any other information that might be important to an object's value.

[0097] Valuation feedback data table **804** includes a user_fdbk_record_id field **826**, an object_id field **828**, a user_id field **830**, a value_category field **832**, a value_feedback field **834**, and an other field **836**. A record is created in valuation feedback data table **804** each time a user provides feedback about the object values associated with a particular data object. The records are generated by policy/value reviser **308** in response to a user providing the feedback.

[0098] User_fdbk_record_id field **826** is the key field of valuation feedback data table **804** and an alphanumeric string, which uniquely identifies each record in valuation feedback data table **804**. Object_id field **828** contains data indicative of a particular data object. Object_id field **828** and object_id field **662** of object metadata table **656** create a many-to-one relationship between records in valuation feedback data table **804** and object metadata table **656**, because each valuation feedback record corresponds to only one object and each object corresponds to many valuation feedback records. User_id field **830** contains data indicative of a particular user. User_id field **830** is analogous to user_id field **632** of user data table **606**. Value_category field **832** contains data indicative of the value category (e.g., continuity value) the feedback is relevant to. Value_feedback field **834** contains data indicative of the feedback regarding the value of the object identified by object_id field **828** in the value category identified by value_category field **832**. For example, value_feedback field **834** might indicate that the object is valued too highly from a continuity perspective. Other field **836** represents one or more additional fields that can contain data indicative of any other information that might be important to a valuation feedback record.

[0099] Currency value data table **806** includes a client_id field **838**, a currency_metric field **840**, a portion_valued_data field **842**, a company_value field **844**, a workforce_access field **846**, a company_revenue field **848**, and an other field **850**. A record is created in currency value data table **806** for each cloud customer. File value generator **306** utilizes currency value data table **806** to

express object values in the form of a monetary currency, such as U.S. dollars.

[0100] Client_id field **838** is the key field of currency value data table **806** and an alphanumeric string, which uniquely identifies a particular cloud customer. Client_id field **838** and client_id field **664** of object metadata table **656** create a one-to-many relationship between records in currency value data table **806** and object metadata table **656**, because each currency value data record corresponds to only one client and each client corresponds to many objects. Through object metadata table **656**, currency value data table **806** has a relationship to many other tables, such as object values table **802** and organizational data table **604**. Currency_metric field **840** contains data indicative of a relationship between a monetary currency and an object value. For example, currency_metric field **840** might indicate that a given object value corresponds to a value of one U.S. dollar. Portion_value_data field **842** contains data indicative of the portion of the cloud customer's file system that is available to be valued. Company_value field **844** contains data indicative of the value (e.g., 10 million dollars) of the cloud customer in a particular currency. Workforce_access field **846** contains data indicative of the portion of the cloud customer's workforce that has access to the files being valued. Company_revenue field **848** contains data indicative of the revenue of the cloud customer (e.g. 1 million dollars per year) in a particular currency. Other field **850** represents one or more additional fields that can contain data indicative of any other information that might be relevant to a currency value data record.

[0101] FIG. **9** shows a database schema **900**, including user data table **606**, object metadata table **656**, and a governance policy data table **902**. It should be noted that, although user data table **606** and object metadata table **656** are shown in more than one of raw data database **404**, compound data database **406**, file value database **408**, and governance policy database **410** they are only saved once in governance database **206** (aside from any backup copies).

[0102] Governance policy data table **902** includes a policy_entry_id field **904**, an object_id field **906**, a user_id field **908**, a backup_no field **910**, a device_type field **912**, a security_req field **914**, and an other field **916**. Policy/value reviser **308** utilizes data from databases **404**, **406**, and **408** to generate governance policy data table **902**.

[0103] Policy_entry_id field **904** is the key field of governance policy data table **902** and an alphanumeric string, which uniquely identifies each record in governance policy data table **902**. Object_id field **906** contains data that uniquely identifies a particular data object. Object_id field **906** and object_id field **662** of object metadata table **656** create a many-to-one relationship between records in governance policy data table **902** and object metadata table **656**, because each governance policy record corresponds to only one object and each object corresponds to many governance policy records. User_id field **908** contains data that uniquely identifies a particular user. User_id field **908** and user_id field **632** of user data table **606** create a many-to-one relationship between records in governance policy data table **902** and user data table **606**, because each governance policy record corresponds to only one user and each user corresponds to many governance policy records. Each record in governance policy data table **902** can include a null value in either of object_id field **906**, user_id field **908**, or neither. In the case of a null value in object_id field **906**, the record corresponds to a particular user without reference to any particular data object. For example, a governance policy record corresponding to a particular user might dictate that the user only access objects having a particular value in one of value categories **522**. In the case of a null value in user_id field **908**, the record corresponds to a particular data object without reference to any particular user. For example, a governance policy record corresponding to a particular data object might dictate that the data object should be password protected. In the case of a null value in neither of object_id field **906** or user_id field **908**, the record corresponds to a particular object with reference to a particular user. For example, a governance policy record corresponding to a particular user and a particular object might dictate that the user access the object from a particular device. Backup_no field **910** contains data indicating the number of times an object should be backed up. Device_type field **912** contains data that indicates what type of

device an object should be stored on. Security_req field **914** contains data that indicates a particular security requirement relating to the object or the user. Other field **916** represents one or more additional fields that can contain data indicative of any other information that might be relevant to a governance policy record.

[0104] FIG. **10** is a flow chart summarizing an example method **1000** for generating object values corresponding to file system objects in a file storage system. In a first step **1002**, secondary data associated with the file system objects is accumulated. The secondary data is data other than (i.e., does not include) content of the file system objects. Then, in a second step **1004**, primary data, including file content, associated with the file system objects is accumulated. Then, in a third step **1006**, the primary data and the secondary data are arranged in a storage device. Next, in a fourth step **1008**, one or more object values indicative of the value(s) of at least one of the file system objects to a particular entity is generated. Finally, in a fifth step **1010**, a record is created in a valuation database. The record associates the object value(s) with the at least one of the file system objects.

[0105] The description of particular embodiments of the present invention is now complete. Many of the described features may be substituted, altered or omitted without departing from the scope of the invention. For example, many, if not all, of the functions described with respect to the remote cloud can also be performed in the local cloud or the valuation service SaaS. As another example, particular data structures are shown to provide examples of data that might be used in the process of generating object values. These data structures are non-limiting examples and can be replaced/supplemented by any of a wide range of varying data structures now known, or yet to be developed, without departing from the scope of the invention. As yet another example, alternate storage and processing devices can be used to facilitate implementation of various aspects of the present invention, including, but not limited to, gate arrays, microprocessors, computer-readable media, and/or semiconductor and/or other optical storage devices. These and other deviations from the particular embodiments shown will be apparent to those skilled in the art, particularly in view of the foregoing disclosure.

Claims

1. A method for performing valuations of file system objects in a file storage system, said method comprising: accumulating secondary data associated with said file system objects, said secondary data being other than content of said file system objects; arranging said secondary data in a storage device; generating an object value indicative of the value of at least one of said file system objects to a particular entity; and creating a record in a valuation database, said record associating said object value with said at least one of said file system objects.
2. The method of claim 1, wherein said step of generating an object value includes generating a plurality of object values, each object value indicative of a distinct value of said at least one of said file system objects from a different perspective.
3. The method of claim 2, wherein one of said plurality of object values is indicative of a value of said at least one of said file system objects from a data security perspective.
4. The method of claim 2, wherein one of said plurality of object values is indicative of a value of said at least one of said file system objects from a business continuity perspective.
5. The method of claim 1, wherein said secondary data includes data associated with said particular entity but not directly associated with any particular subset of said file system objects.
6. The method of claim 1, wherein said object value is expressed in the form of a particular monetary currency.
7. The method of claim 1, wherein said step of generating an object value includes: generating compound data based at least in part on said secondary data accumulated from different sources; and using said compound data to generate said object value.

- 8.** The method of claim 7, wherein said compound data includes user activity data associated with said file system objects, said user activity data indicative of activity of users of said file storage system with respect to said file system objects.
- 9.** The method of claim 1, further comprising: accumulating primary data indicative of content of said individual ones of said file system objects; and arranging said primary data in said storage device.
- 10.** The method of claim 9, wherein said step of accumulating primary data includes scanning said individual ones of said file system objects to obtain said content from said individual ones of said file system objects.
- 11.** The method of claim 9, wherein said step of accumulating primary data includes categorizing said individual ones of said file system objects based at least in part on said content of said individual ones of said file system objects.
- 12.** The method of claim 9, further comprising utilizing at least a portion of said primary data and said secondary data to generate compound data, said compound data being indicative of relationships between individual ones of said file system objects.
- 13.** The method of claim 12, wherein said step of utilizing at least a portion of said primary data and said secondary data to generate compound data includes estimating the similarity of two or more different ones of said file system objects, based at least in part on said primary data.
- 14.** The method of claim 1, wherein said step of accumulating secondary data includes accessing metadata associated with said individual ones of said file system objects.
- 15.** The method of claim 1, wherein said step of accumulating secondary data includes accumulating data indicative of an infrastructure of said file storage system.
- 16.** The method of claim 1, further comprising expressing said object value of said at least one of said file system objects relative to object values of others of said file system objects.
- 17.** The method of claim 16, wherein said step of expressing said object value of said at least one of said file system objects relative to object values of others of said file system objects includes expressing said object value relative to a base value that is equal to a lowest object value of said object values of said others of said file system objects.
- 18.** The method of claim 1, further comprising utilizing said object value of said at least one of said file system objects to inform IT policy changes relating to said particular entity.
- 19.** The method of claim 1, wherein said step of accumulating secondary data includes accessing said file system objects of said file storage system over a wide area network.
- 20.** A valuation system for performing valuations of file system objects, said valuation system comprising: a processor configured to execute code; memory configured to store data and said code, said code including an accumulator configured to accumulate secondary data associated with said file system objects, said secondary data being non-object-content data, a data aggregator configured to arrange said secondary data in said memory, and a value generator configured to generate an object value indicative of the value of at least one of said file system objects to a particular entity, and to create a record in a valuation database to associate said object value with said at least one of said file system objects.
- 21-39.** (canceled)
-