



US 20250265772A1

(19) **United States**

(12) **Patent Application Publication**

MANDAL et al.

(10) **Pub. No.: US 2025/0265772 A1**

(43) **Pub. Date:** **Aug. 21, 2025**

(54) **OCCLUSION CULLING FOR DEPTH
VALUES IN A DEPTH BUFFER**

(52) **U.S. Cl.**
CPC **G06T 15/405** (2013.01)

(71) Applicant: **Intel Corporation**, Santa Clara, CA
(US)

(57) **ABSTRACT**

(72) Inventors: **Saikat MANDAL**, Sacramento, CA
(US); **Ranjani NARAYANAN**,
Cupertino, CA (US); **Sunil NATARAJ**,
Milpitas, CA (US)

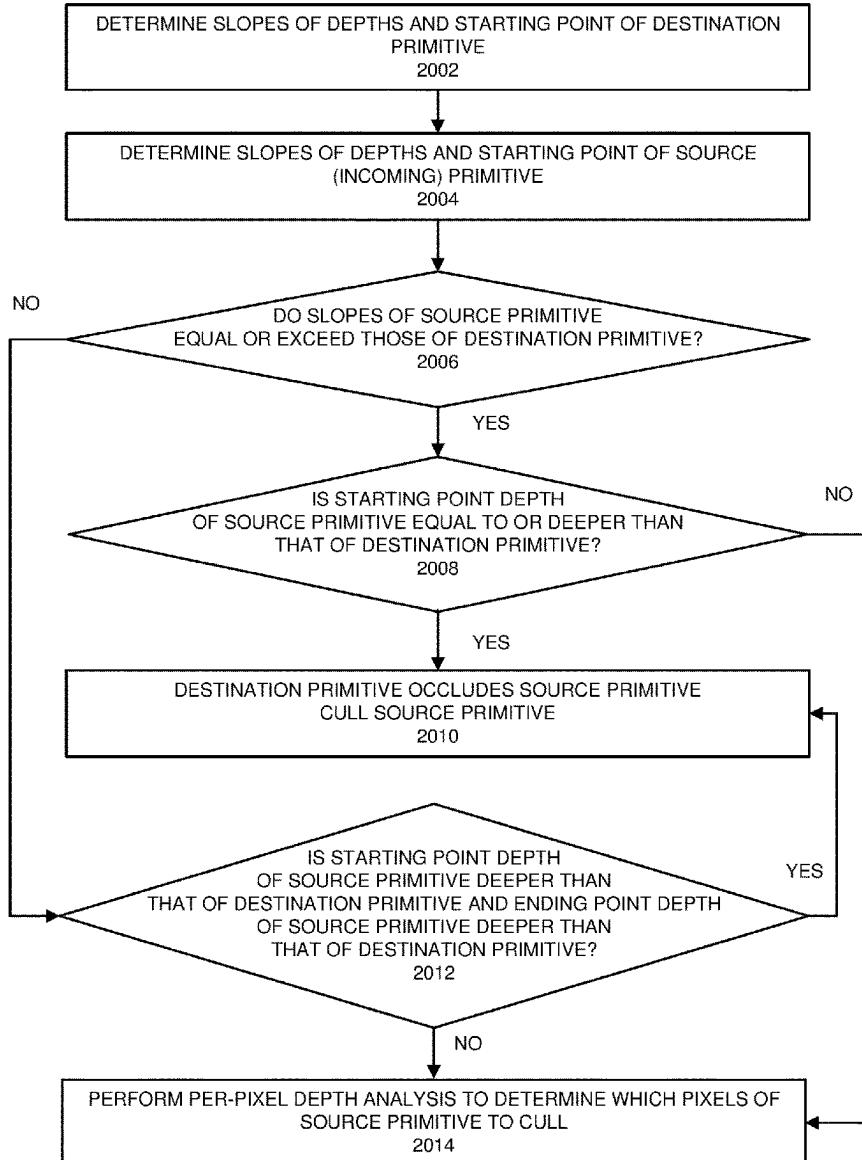
(21) Appl. No.: **18/583,773**

(22) Filed: **Feb. 21, 2024**

Publication Classification

(51) **Int. Cl.**
G06T 15/40 (2011.01)

A system that includes a graphics processing unit (GPU) to determine whether the second primitive is potentially occluded, fully occluded, or not occluded by the first primitive based on a comparison of the minimum depth value of the first primitive and the maximum depth value of the second primitive and a comparison of the depth slope data of the first primitive and the depth slope data of the second primitive. Based on a determination the second primitive is fully occluded by the first primitive, the GPU is to exclude the second primitive from rendering.



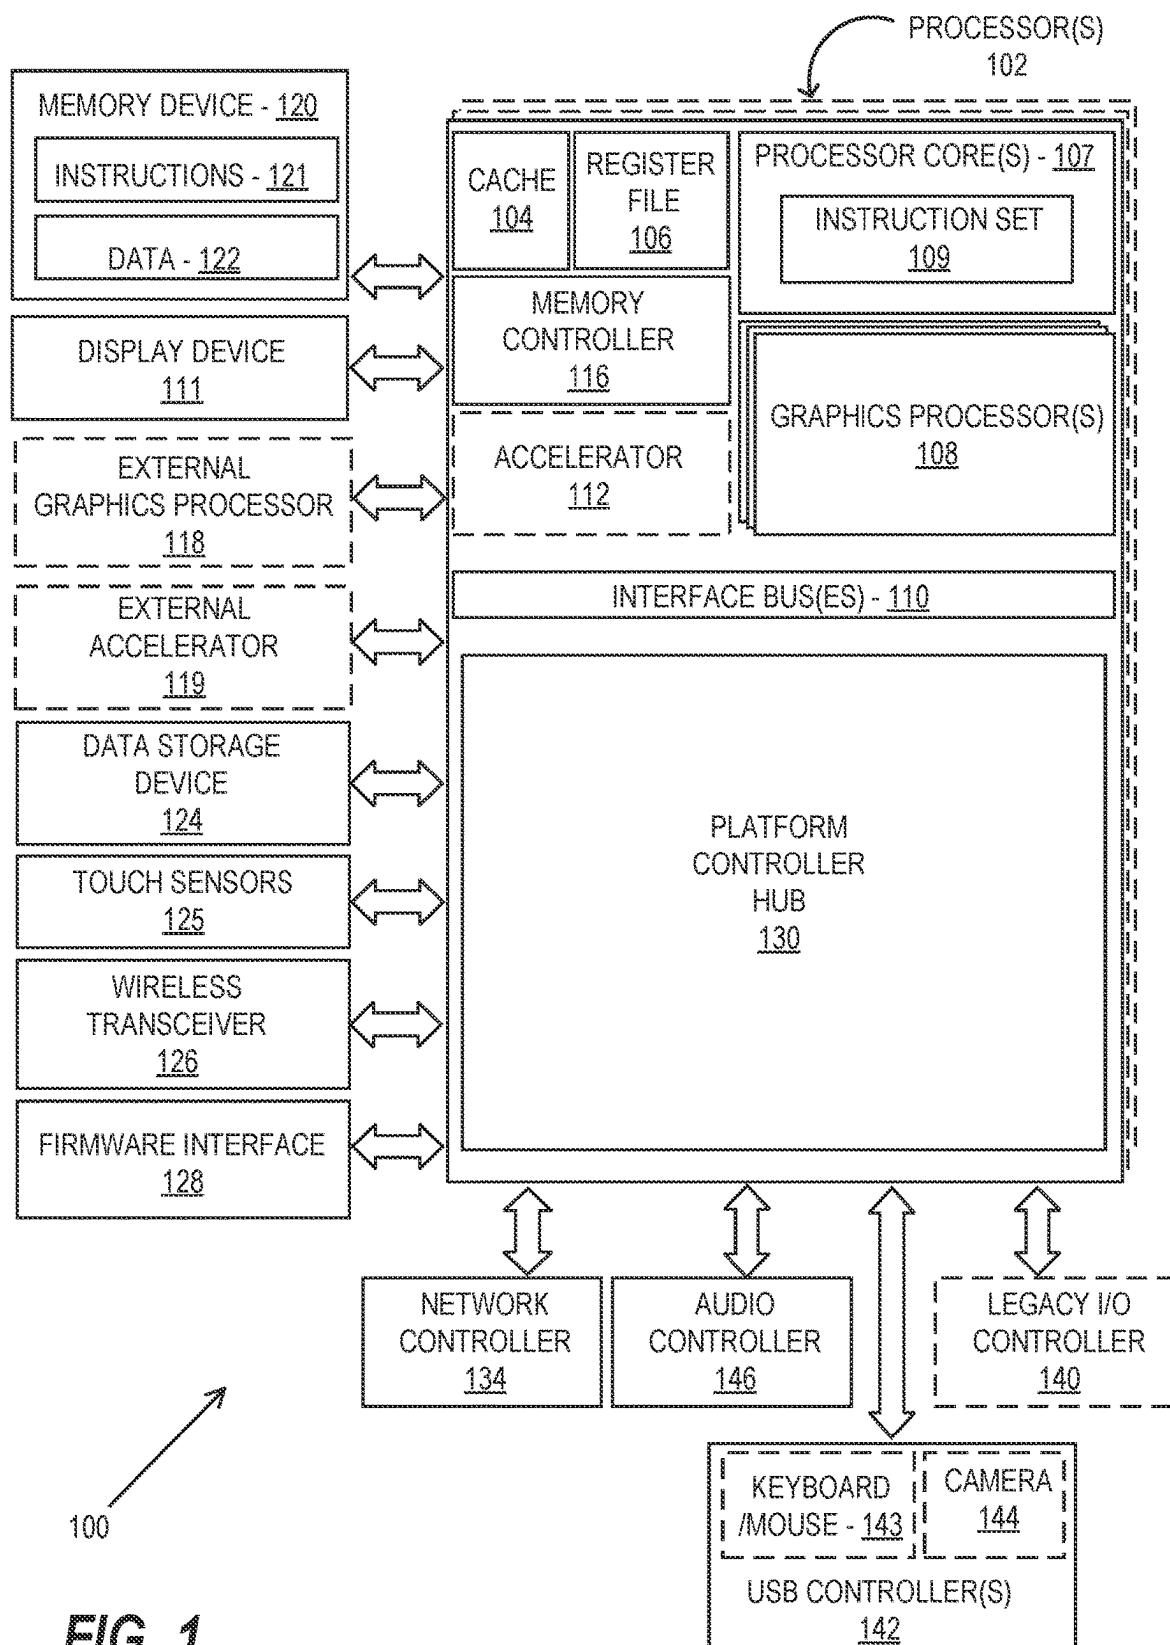


FIG. 1

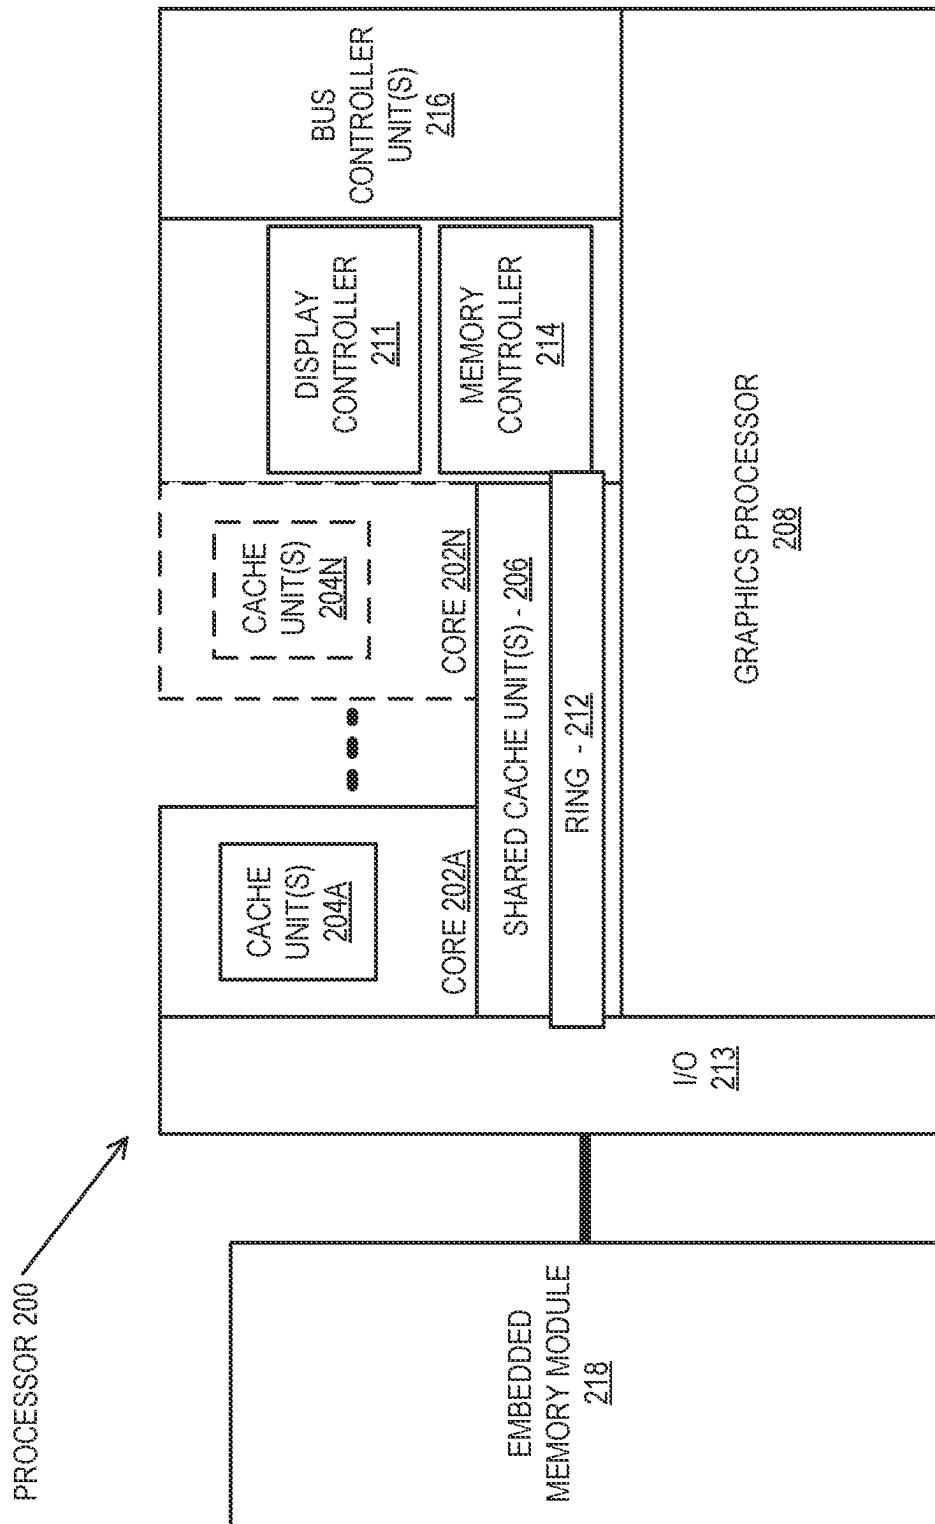


FIG. 2A

219

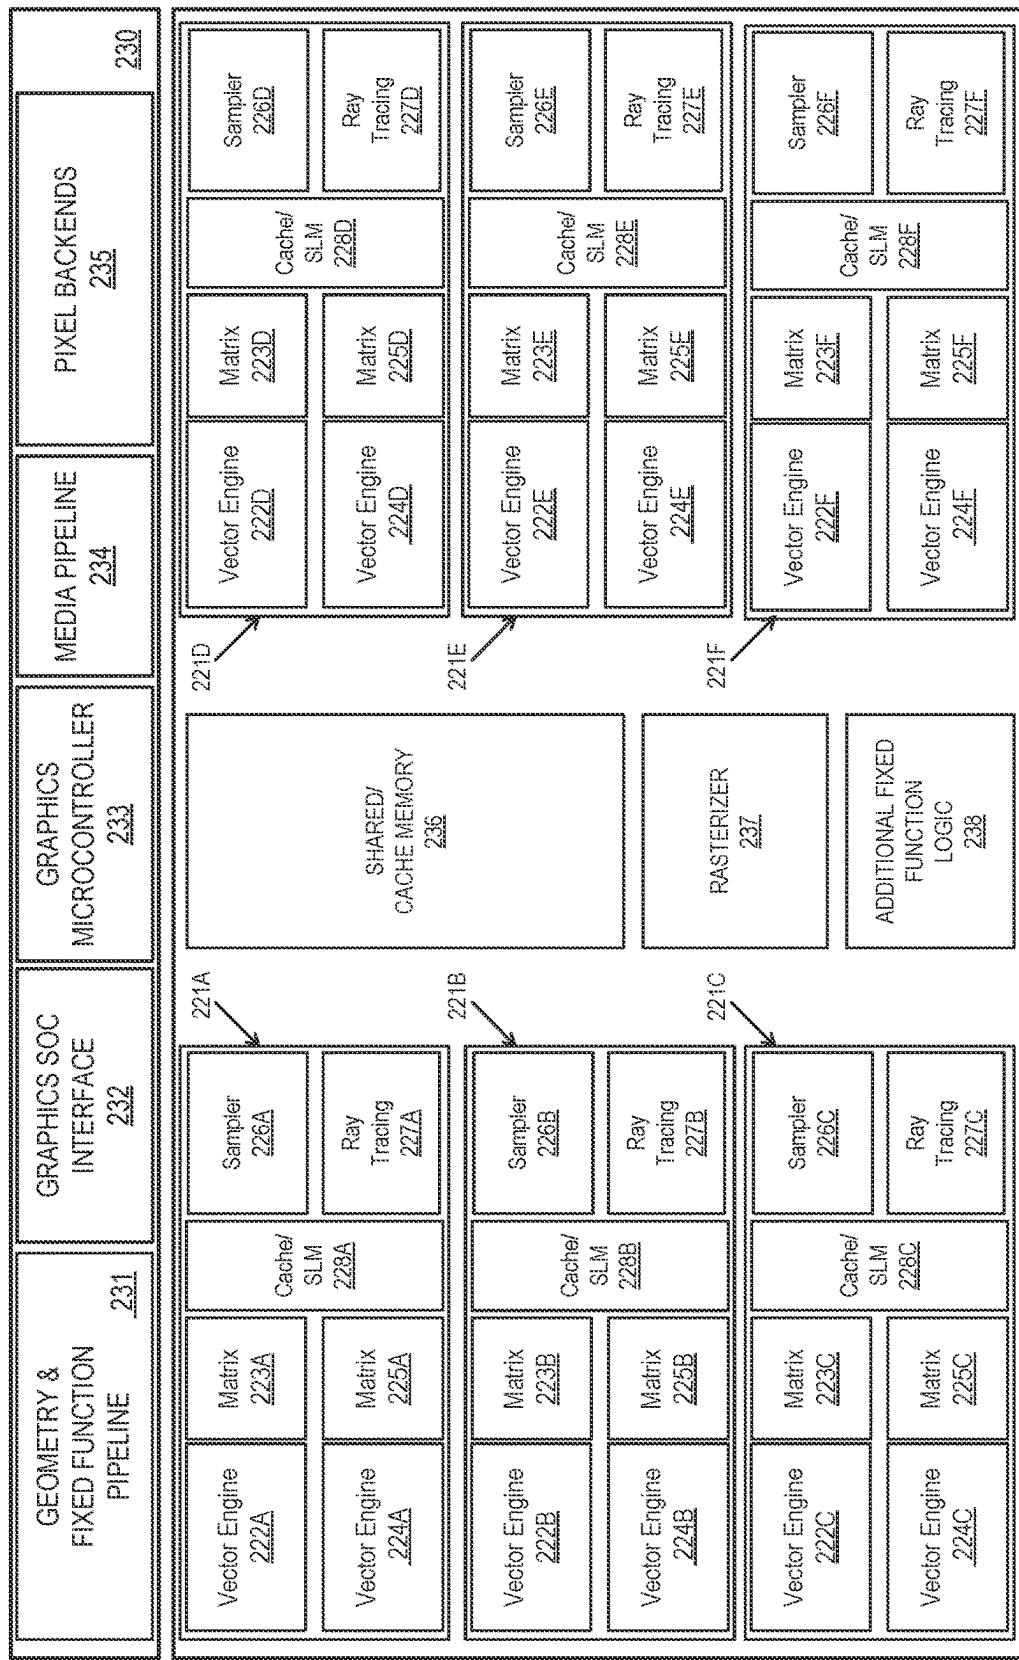


FIG. 2B

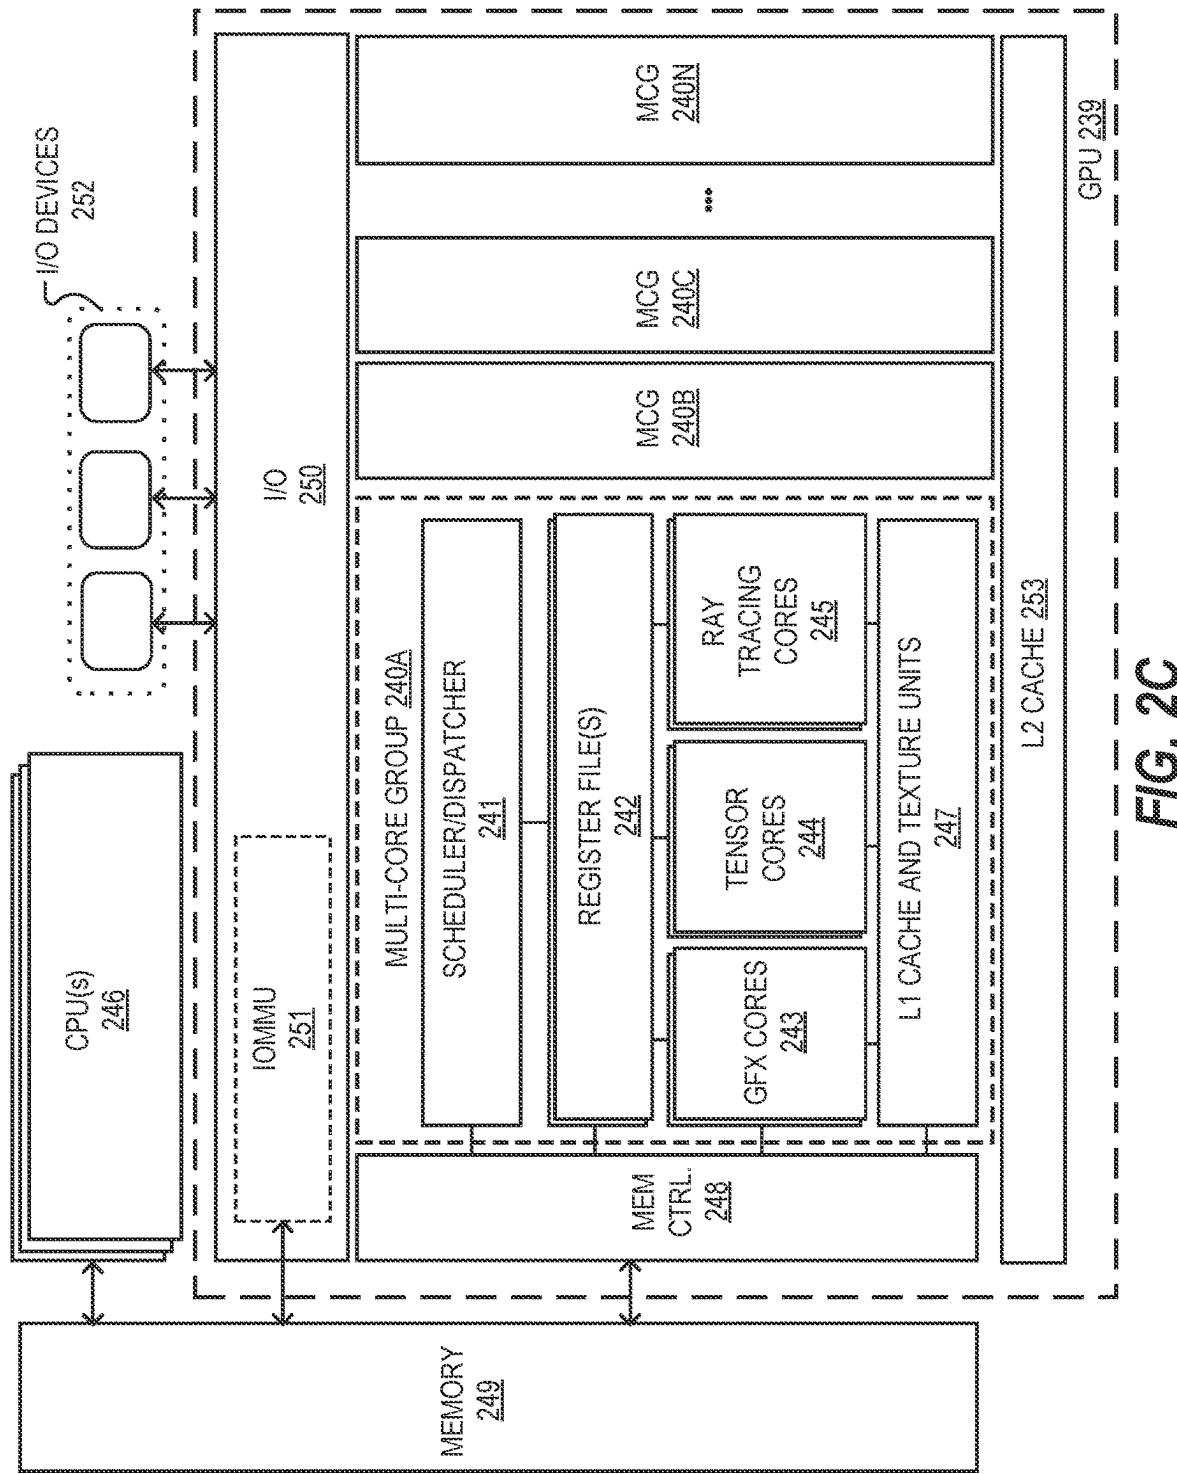


FIG. 2C

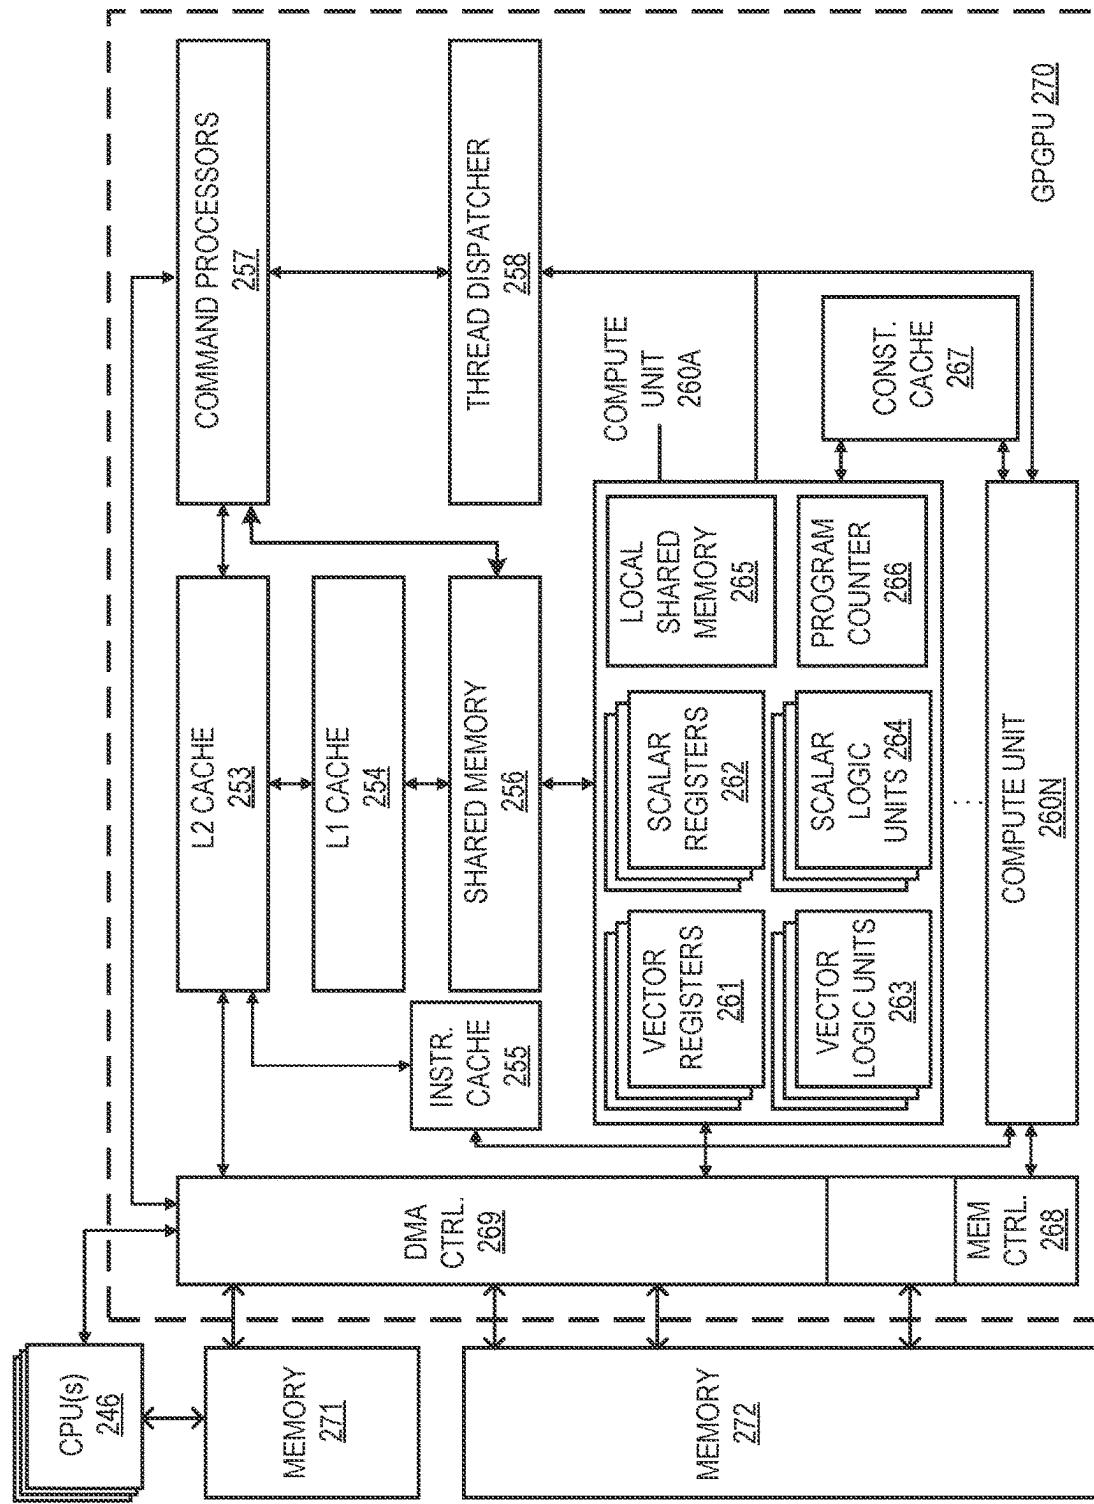


FIG. 2D

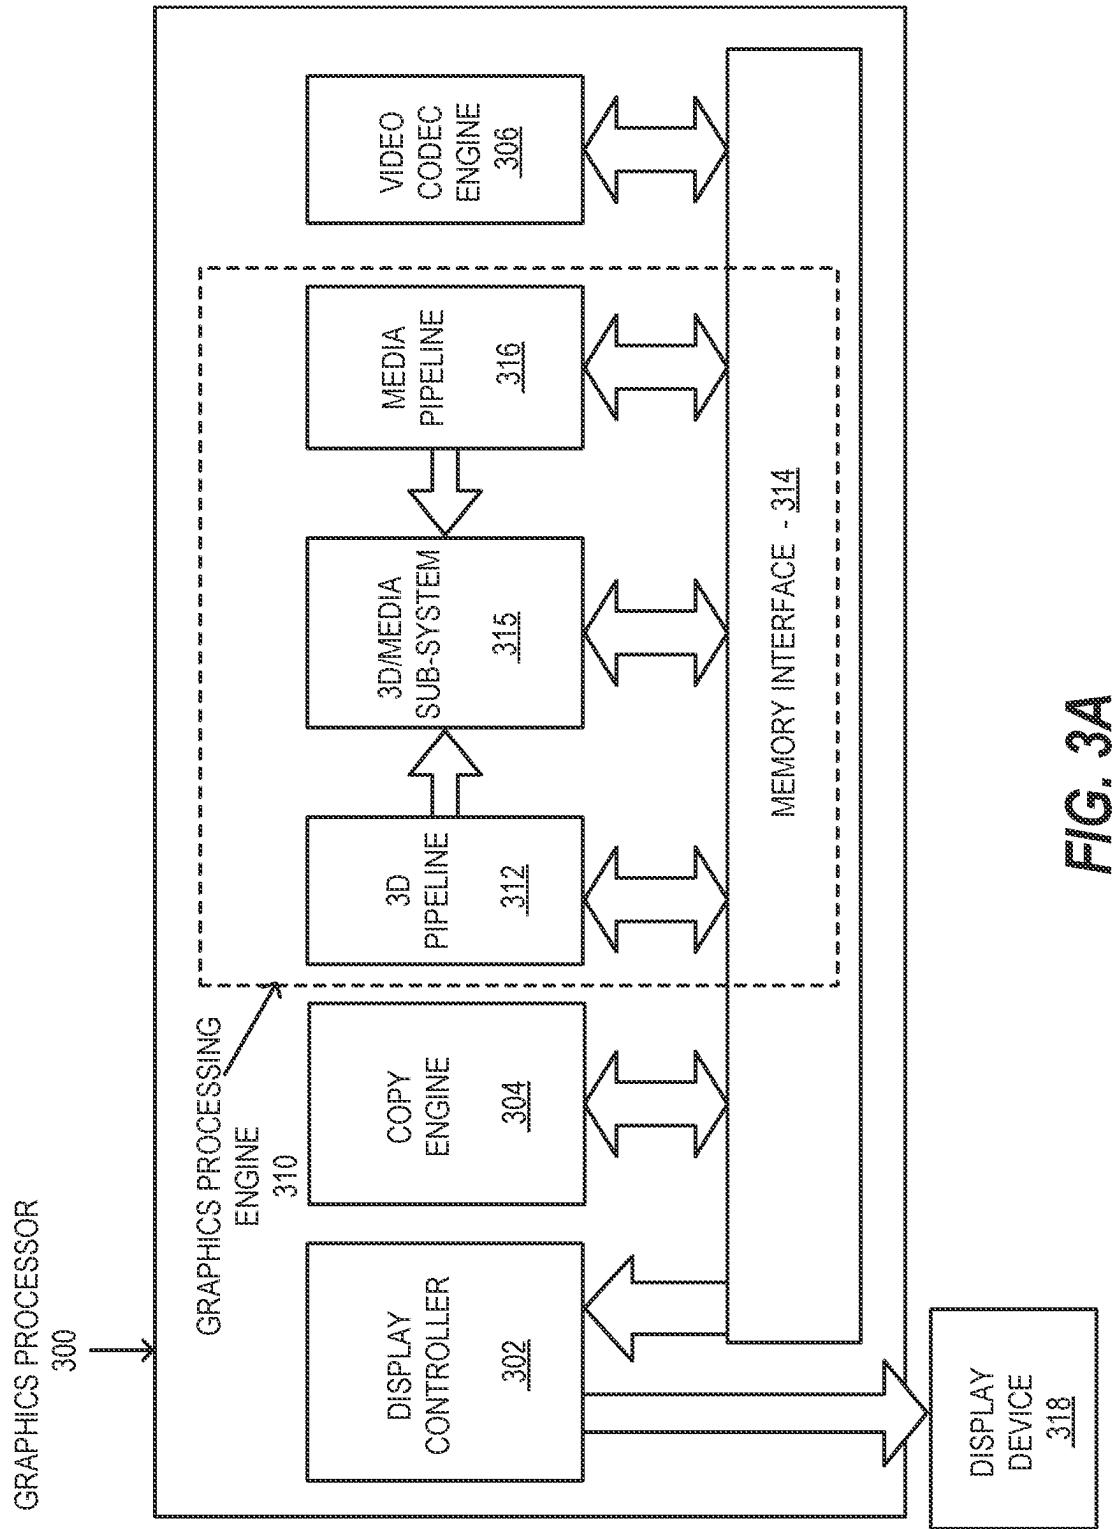


FIG. 3A

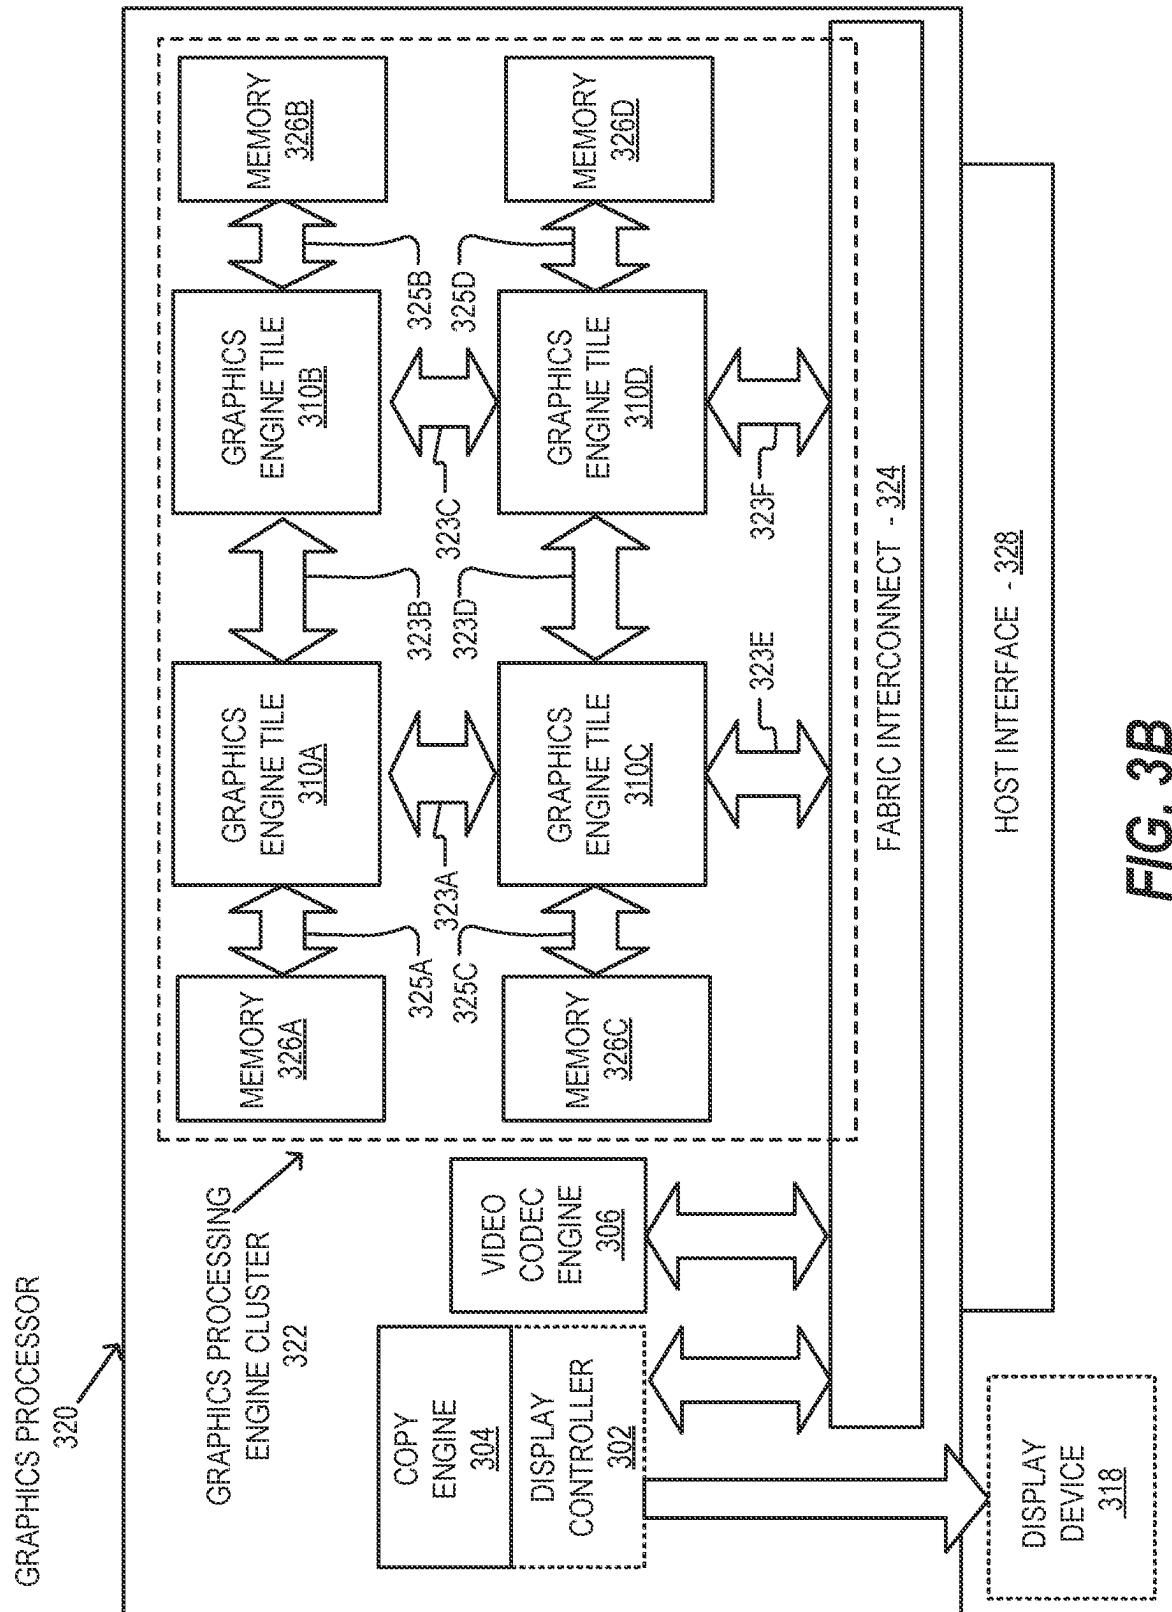


FIG. 3B

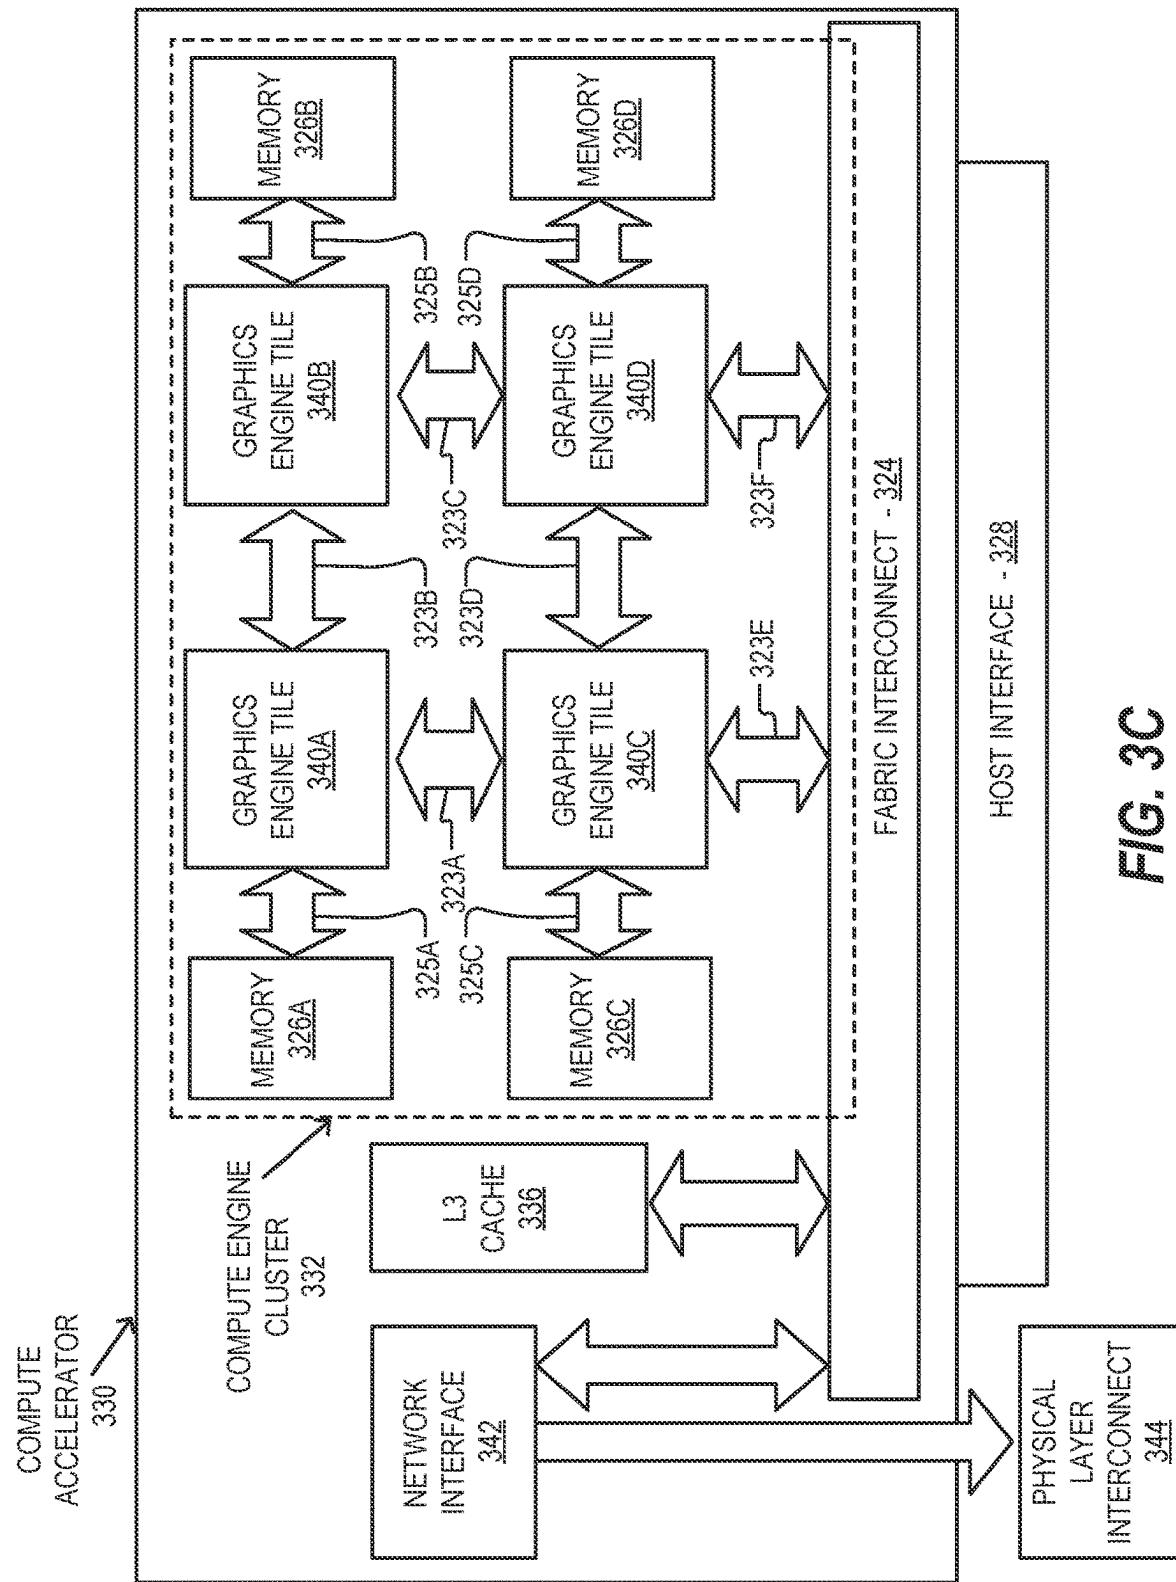


FIG. 3C

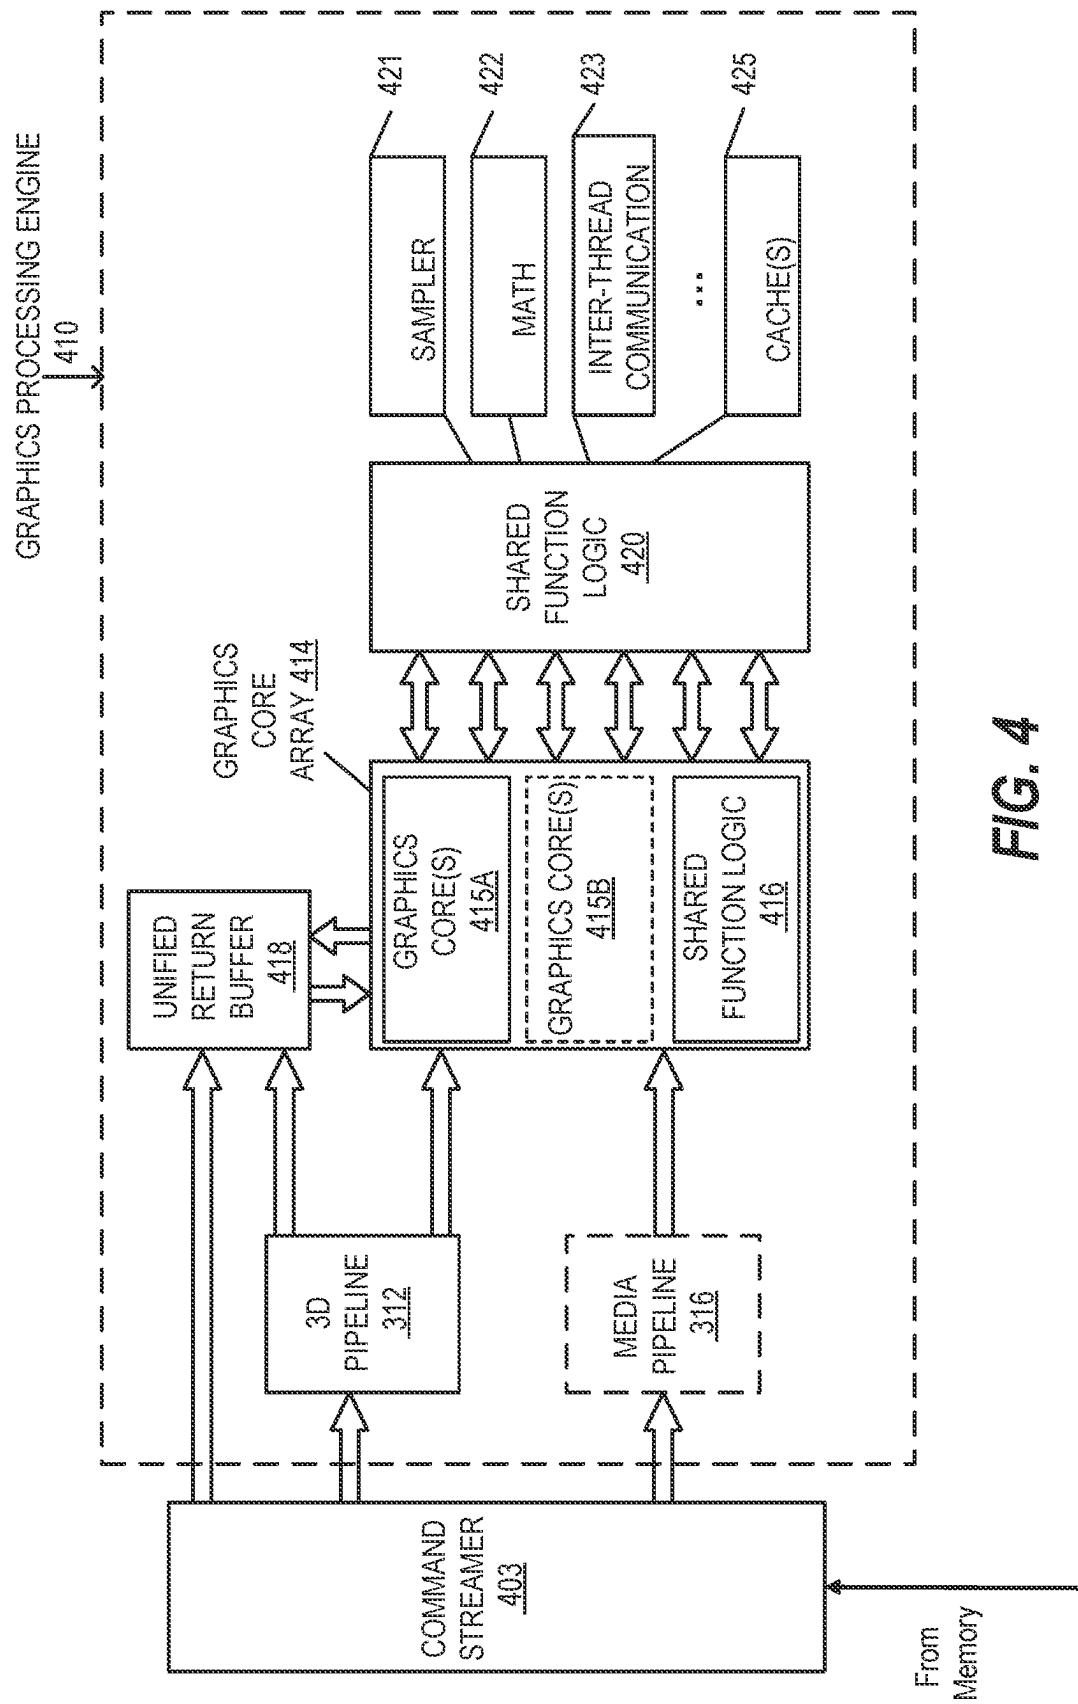


FIG. 4

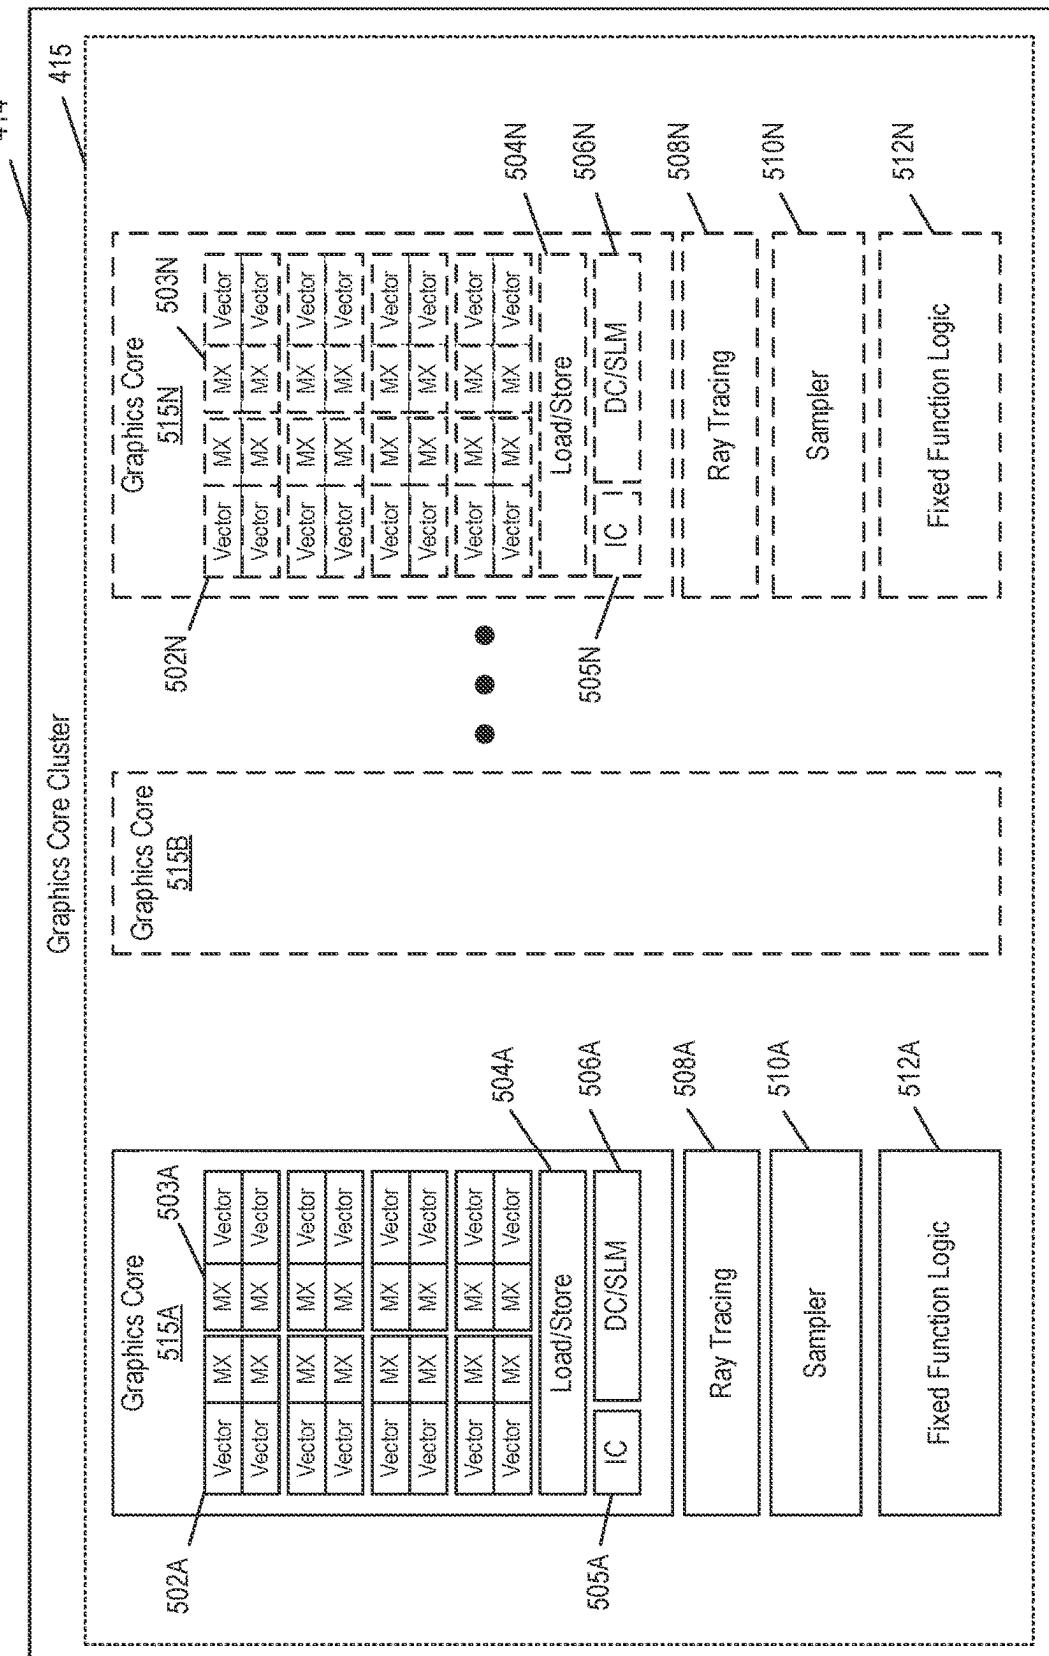


FIG. 5A

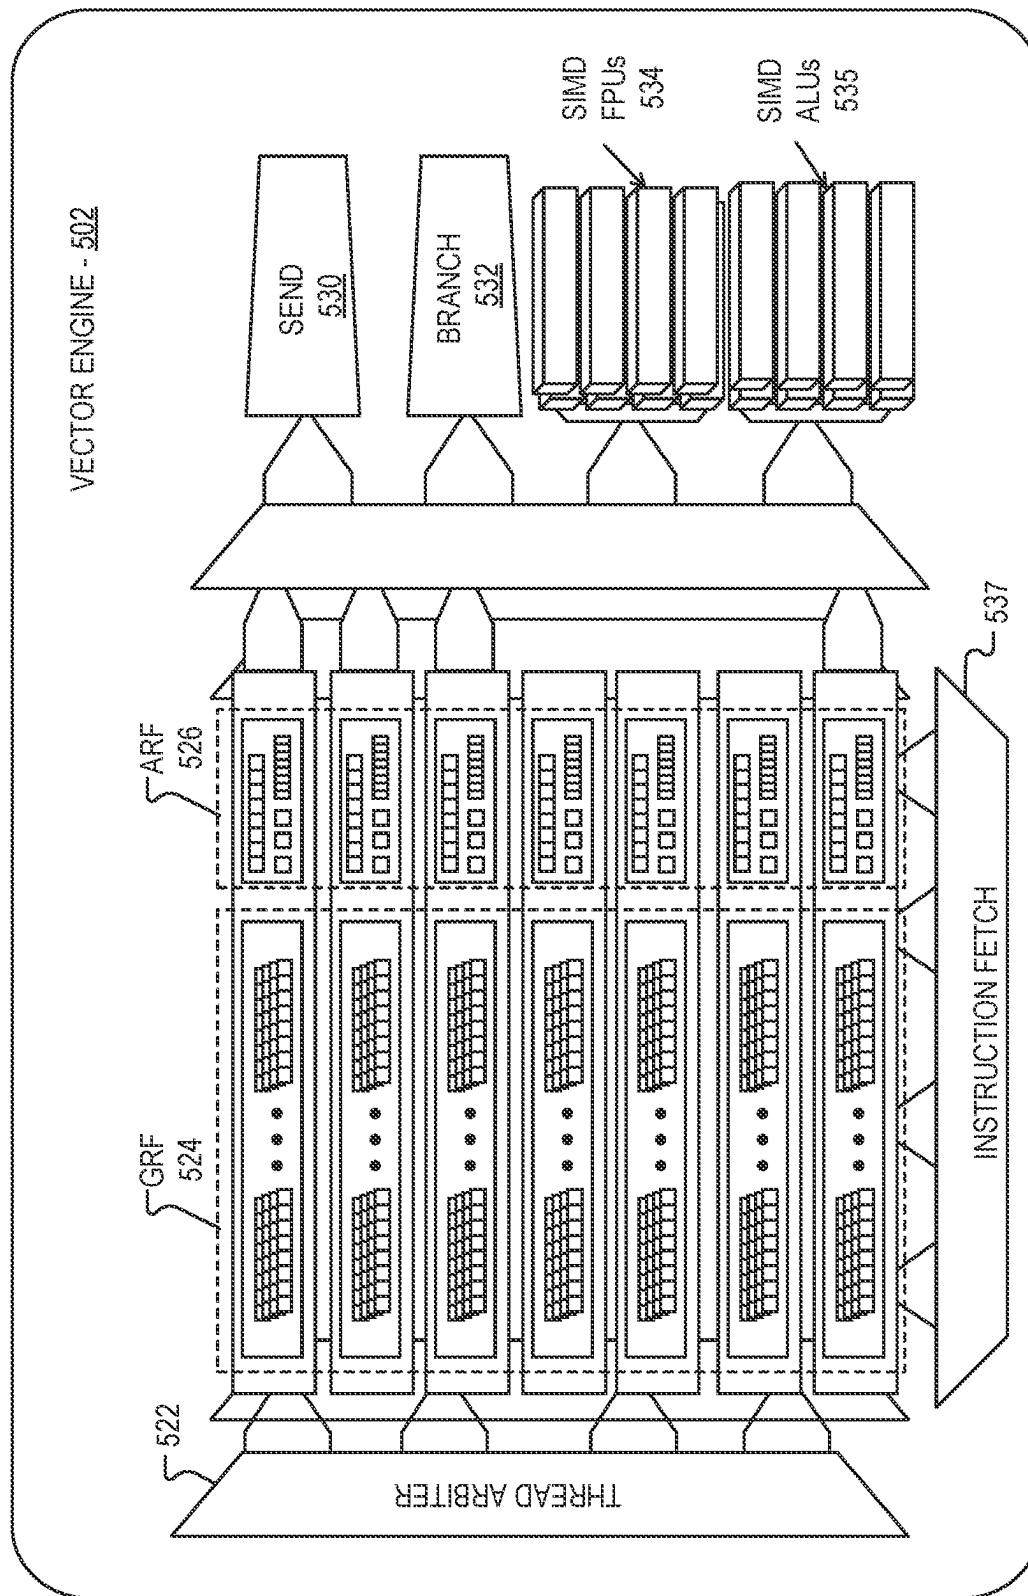


FIG. 5B

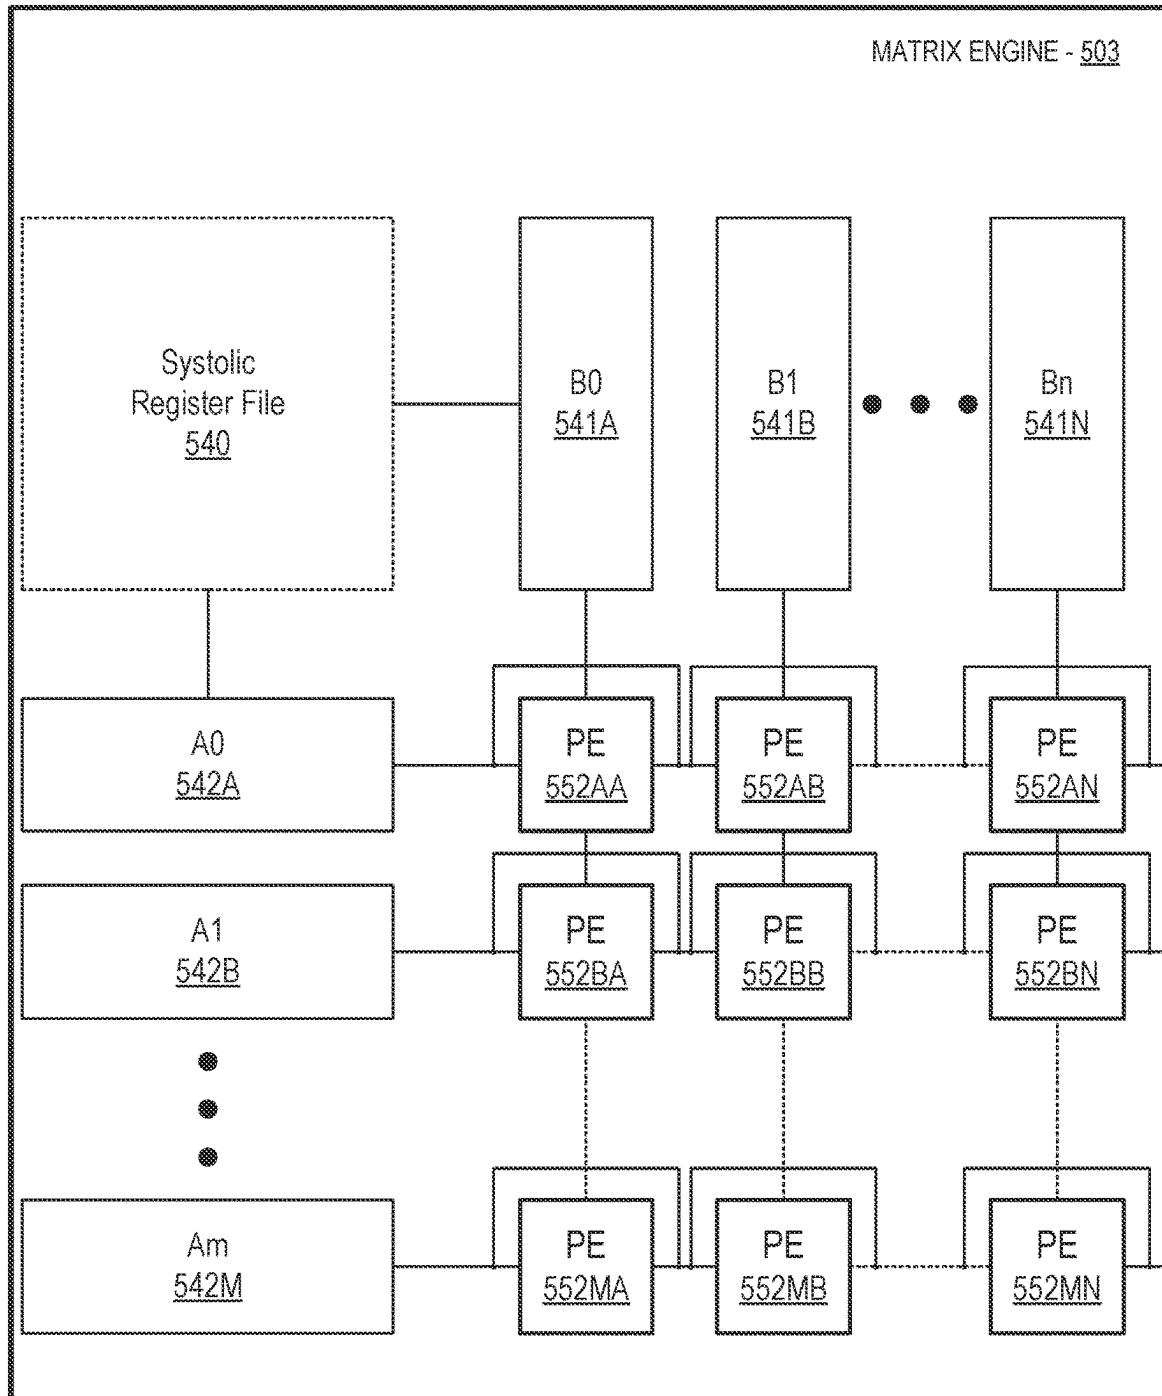
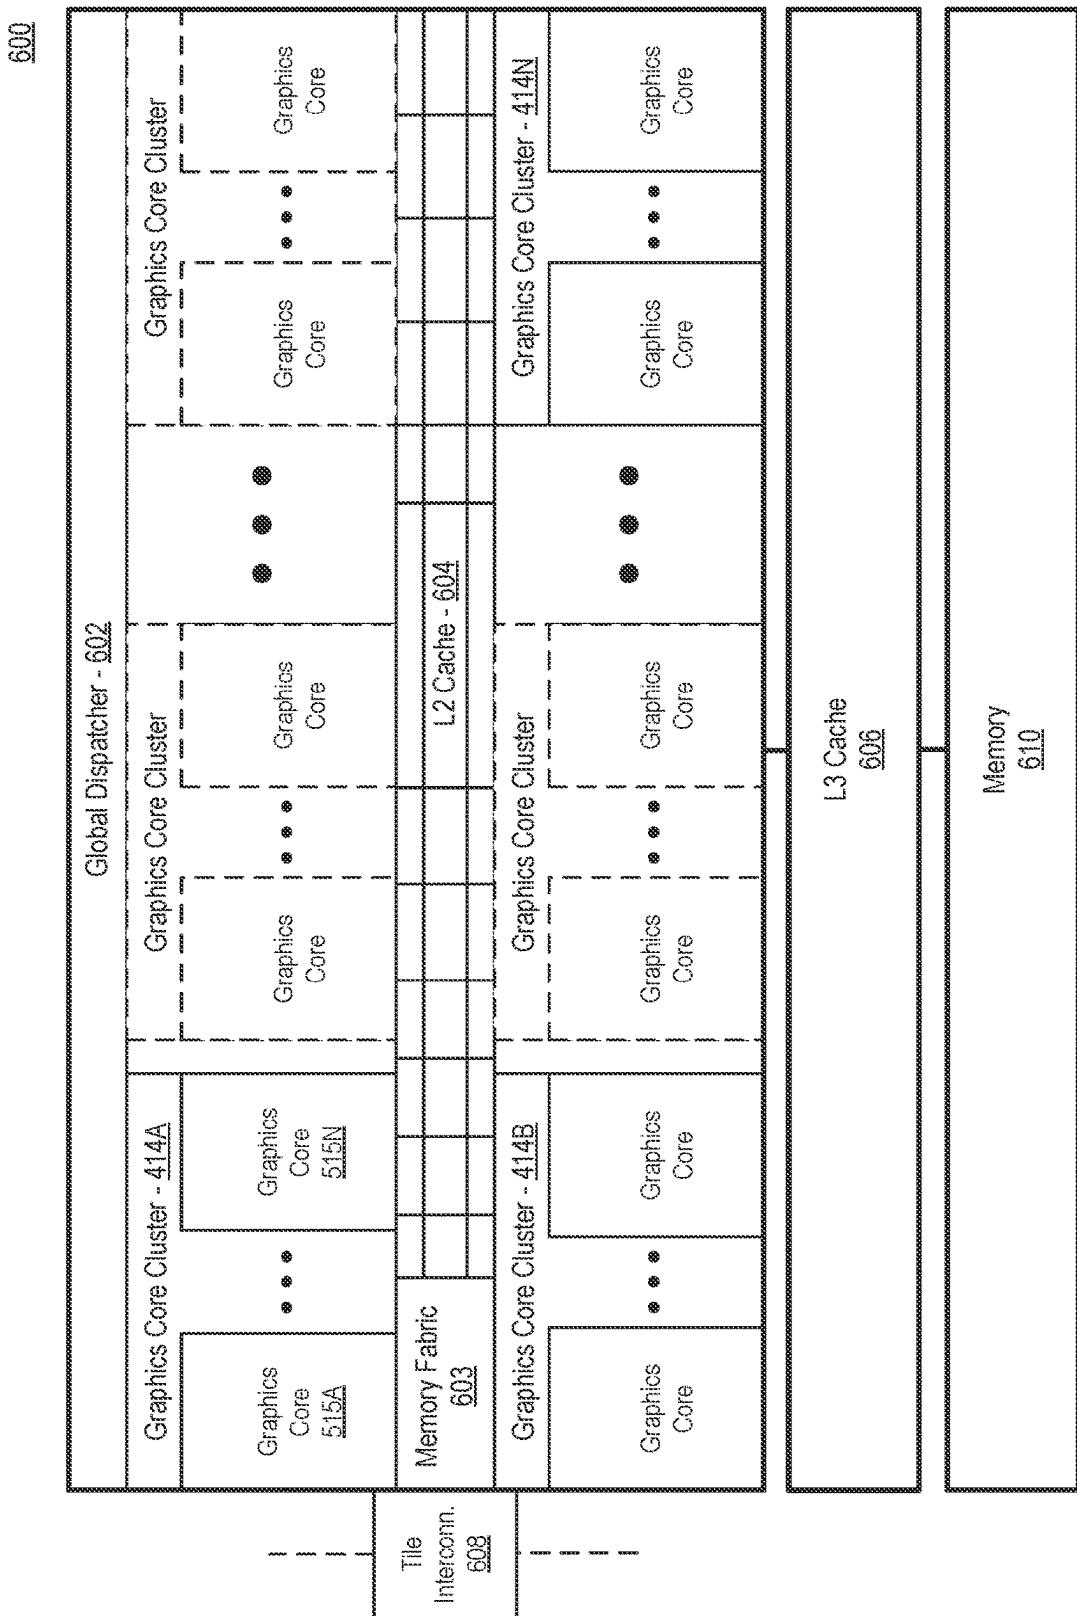


FIG. 5C



6
FIG.^{*}

GRAPHICS PROCESSOR INSTRUCTION FORMATS

700

128-BIT INSTRUCTION

710

OPCODE <u>712</u>	CONTROL <u>714</u>	EXEC-SIZE <u>716</u>	DEST <u>718</u>	SRC0 <u>720</u>	SRC1 <u>722</u>	SRC2 <u>724</u>	ACCESS/ADDRESS MODE <u>726</u>
----------------------	-----------------------	-------------------------	--------------------	--------------------	--------------------	--------------------	-----------------------------------

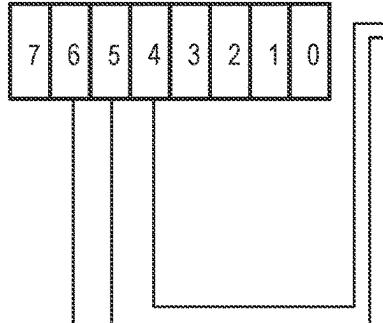
64-BIT COMPACT INSTRUCTION

730

OPCODE <u>712</u>	INDEX <u>713</u>	CONTROL <u>714</u>	DEST <u>718</u>	SRC0 <u>720</u>	SRC1 <u>722</u>
----------------------	---------------------	-----------------------	--------------------	--------------------	--------------------

OPCODE DECODE

740



- opcode=000xxxxb ← Move/Logic - 742
- opcode=0010xxxxb ← Miscellaneous - 746
- opcode=0011xxxxb ← Flow Control - 744
- opcode=0100xxxxb ← Parallel Math - 748
- opcode=0101xxxxb ← Vector Math - 750

FIG. 7

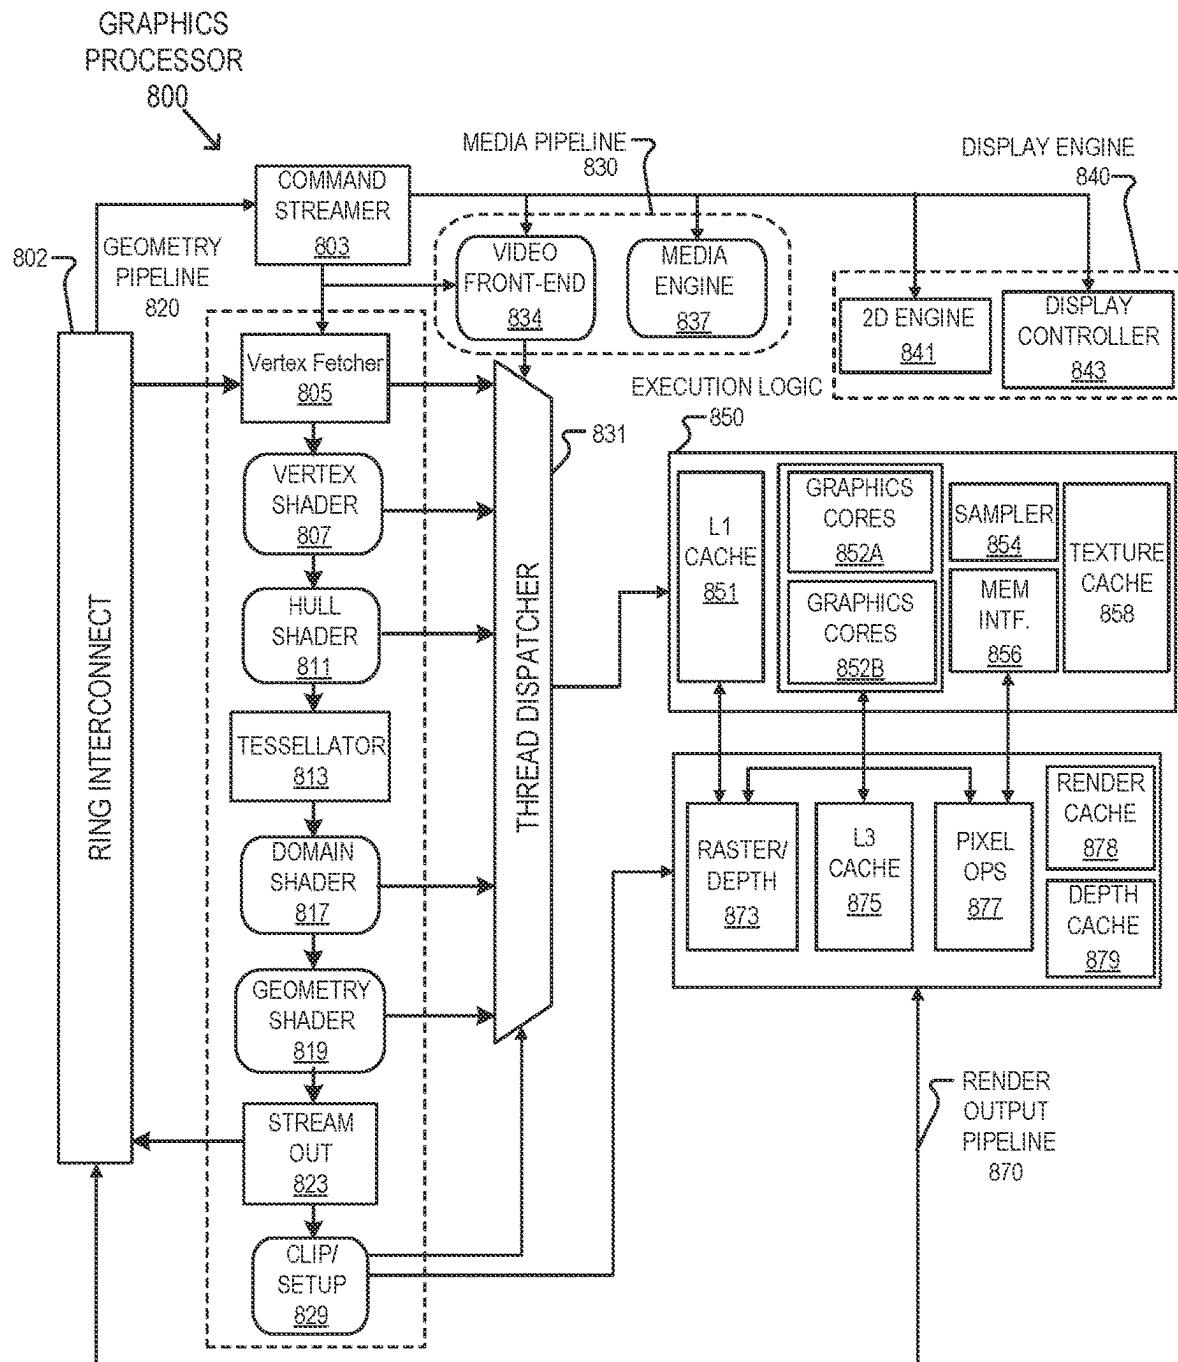


FIG. 8

FIG. 9A

GRAPHICS PROCESSOR COMMAND
FORMAT
900

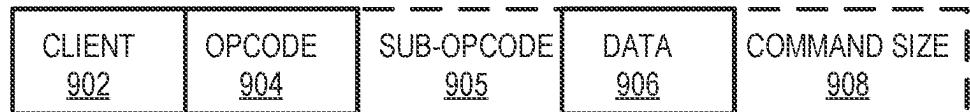
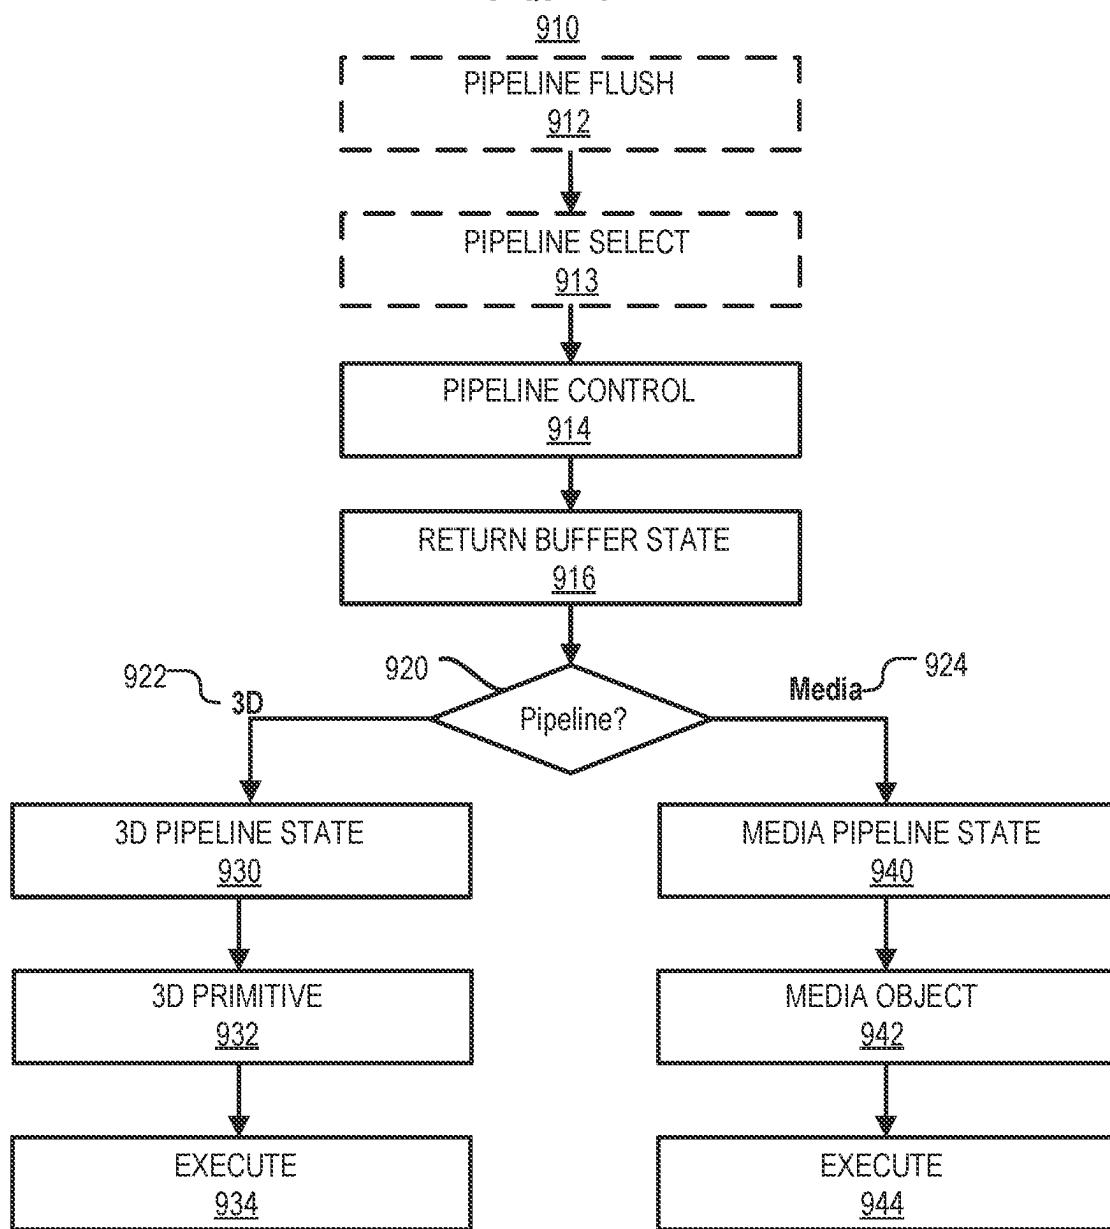


FIG. 9B

GRAPHICS PROCESSOR COMMAND
SEQUENCE



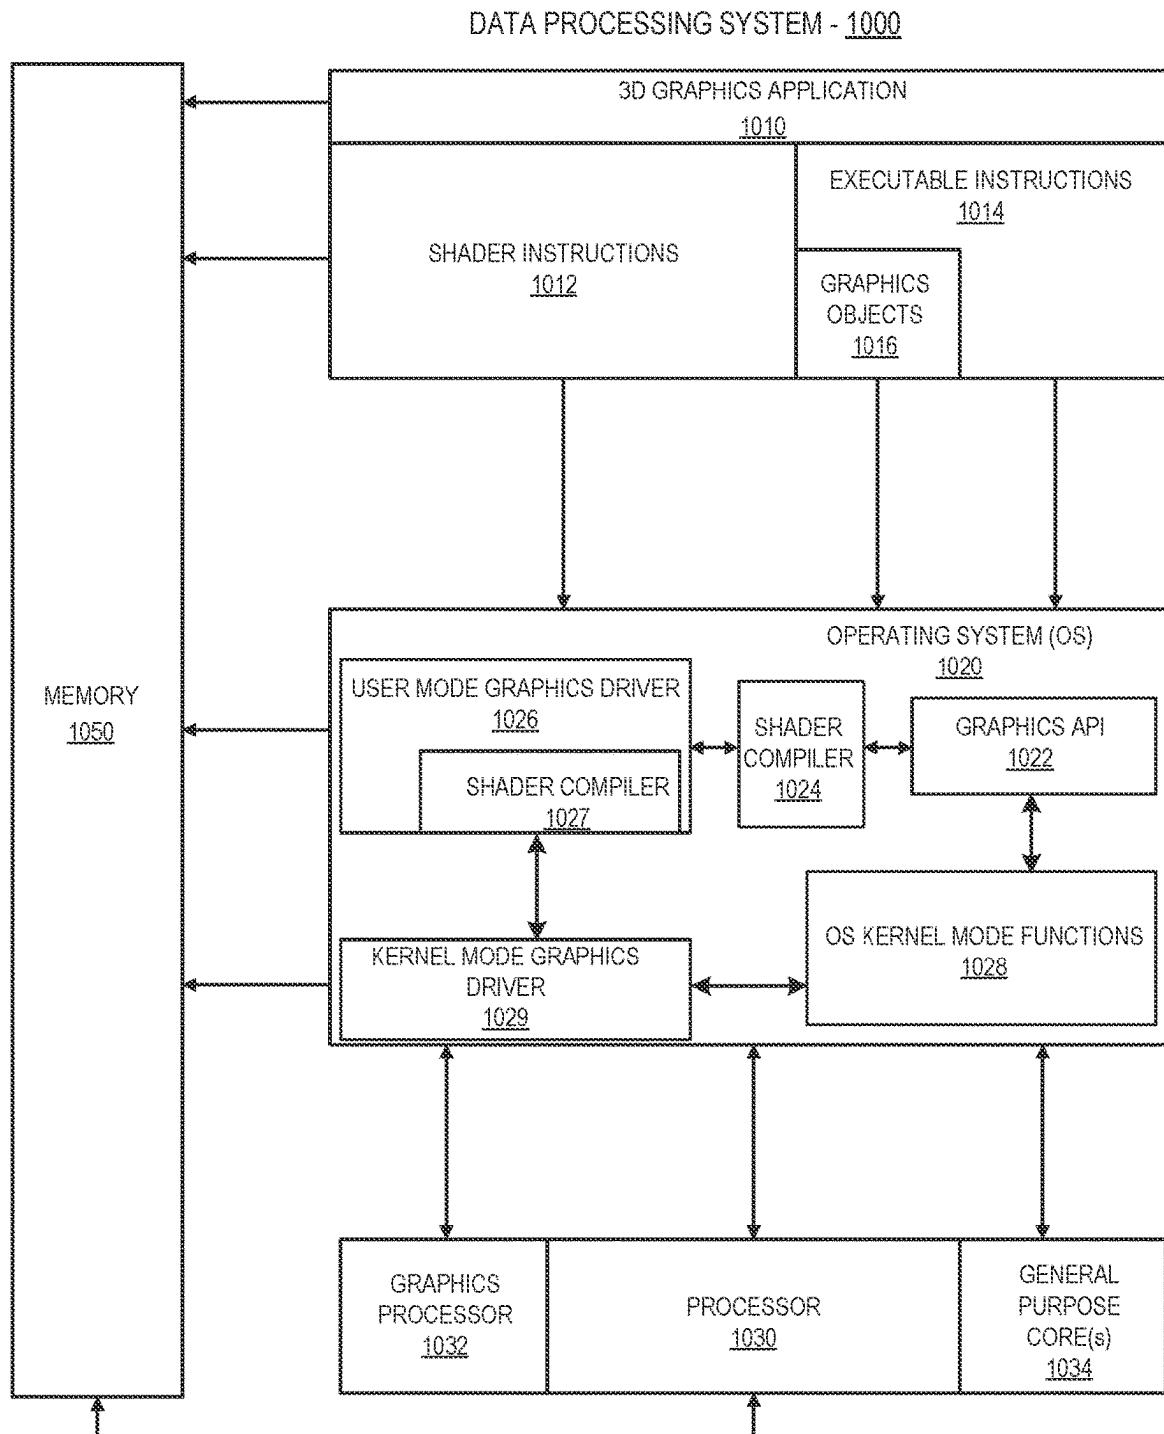


FIG. 10

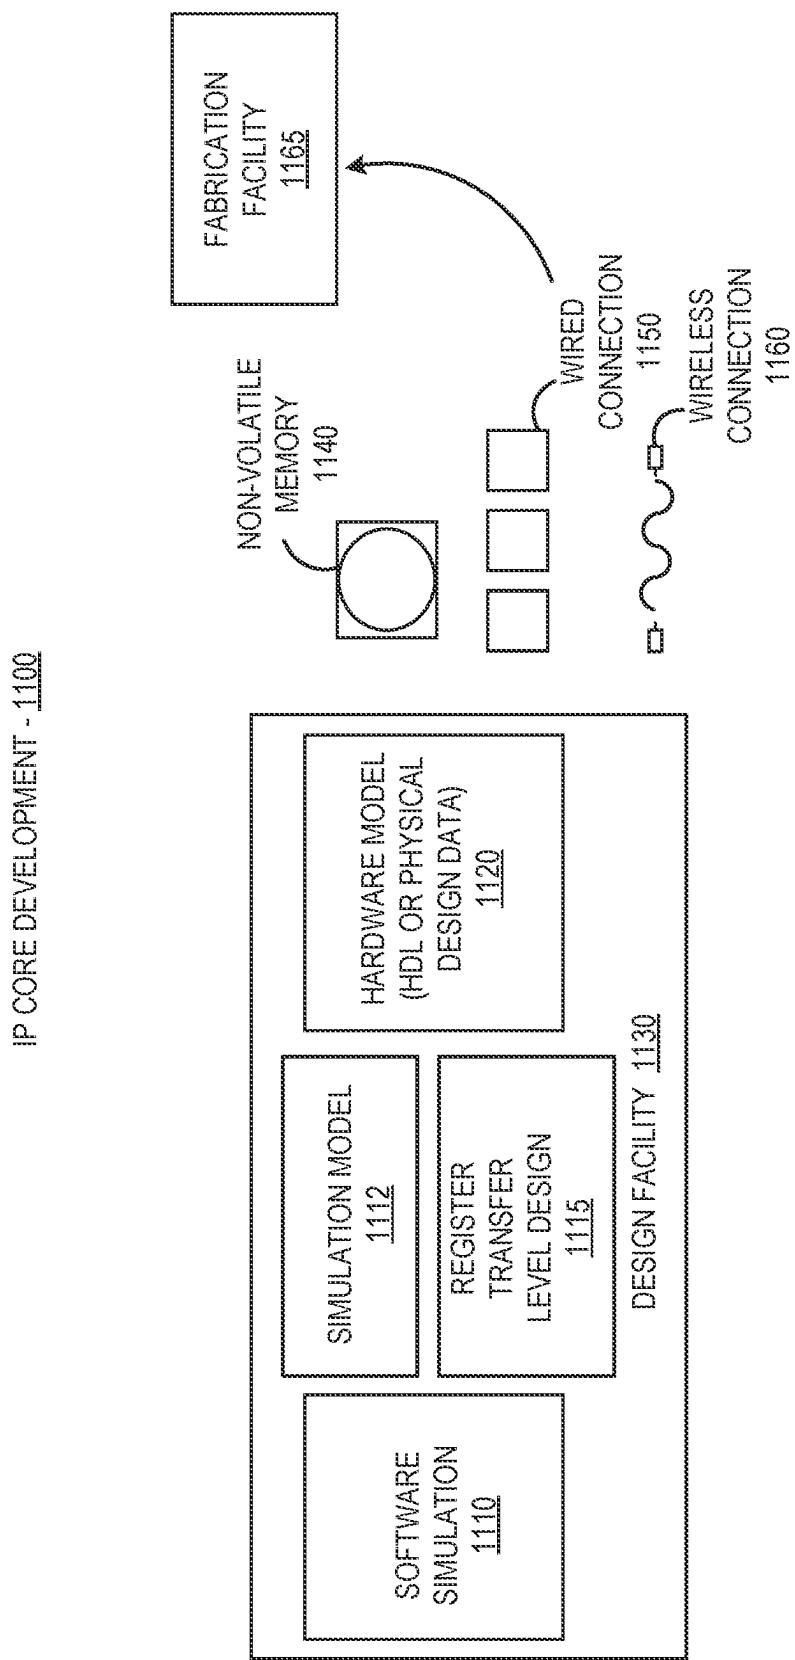


FIG. 11A

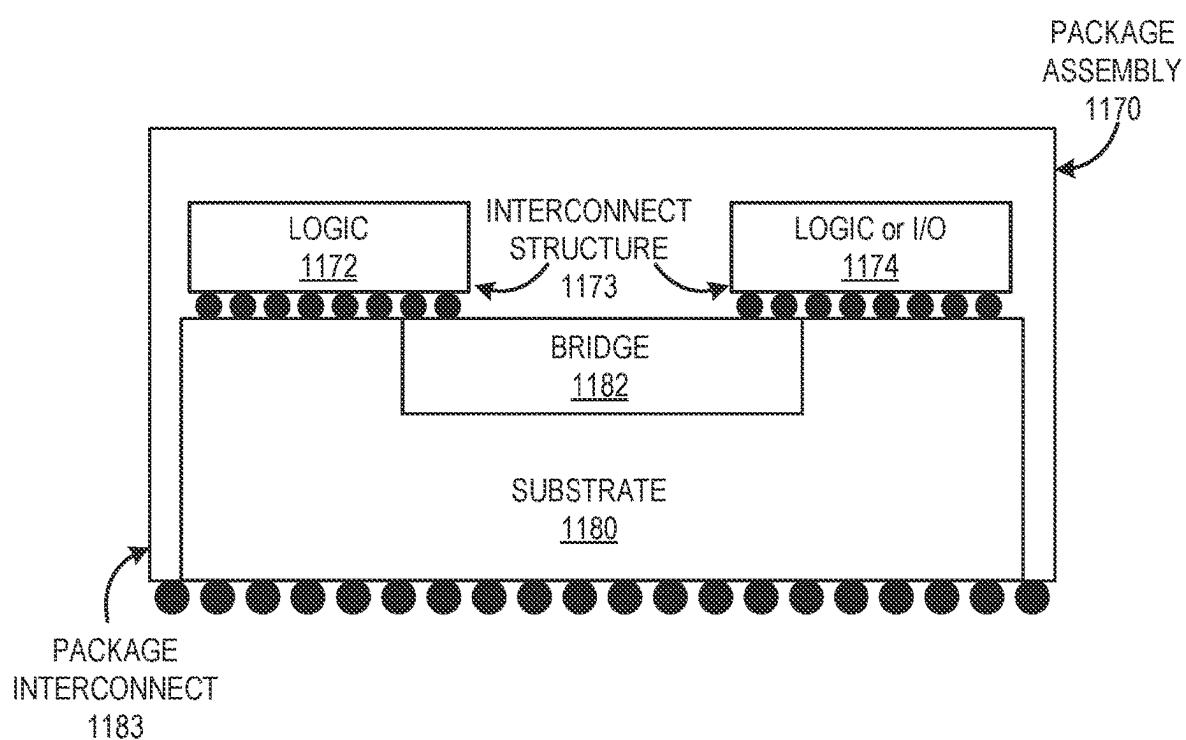


FIG. 11B

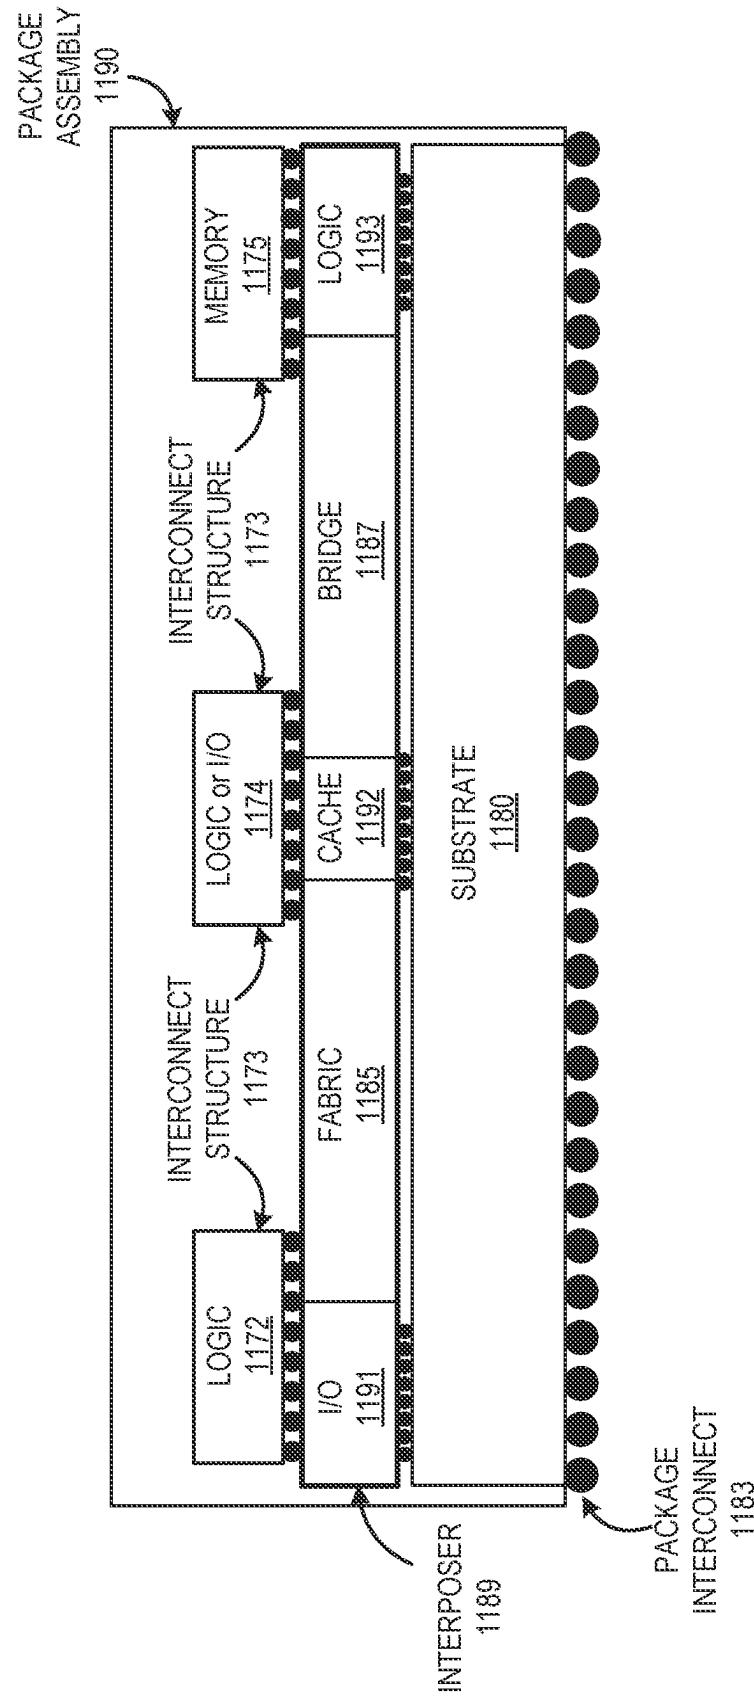


FIG. 11C

1194

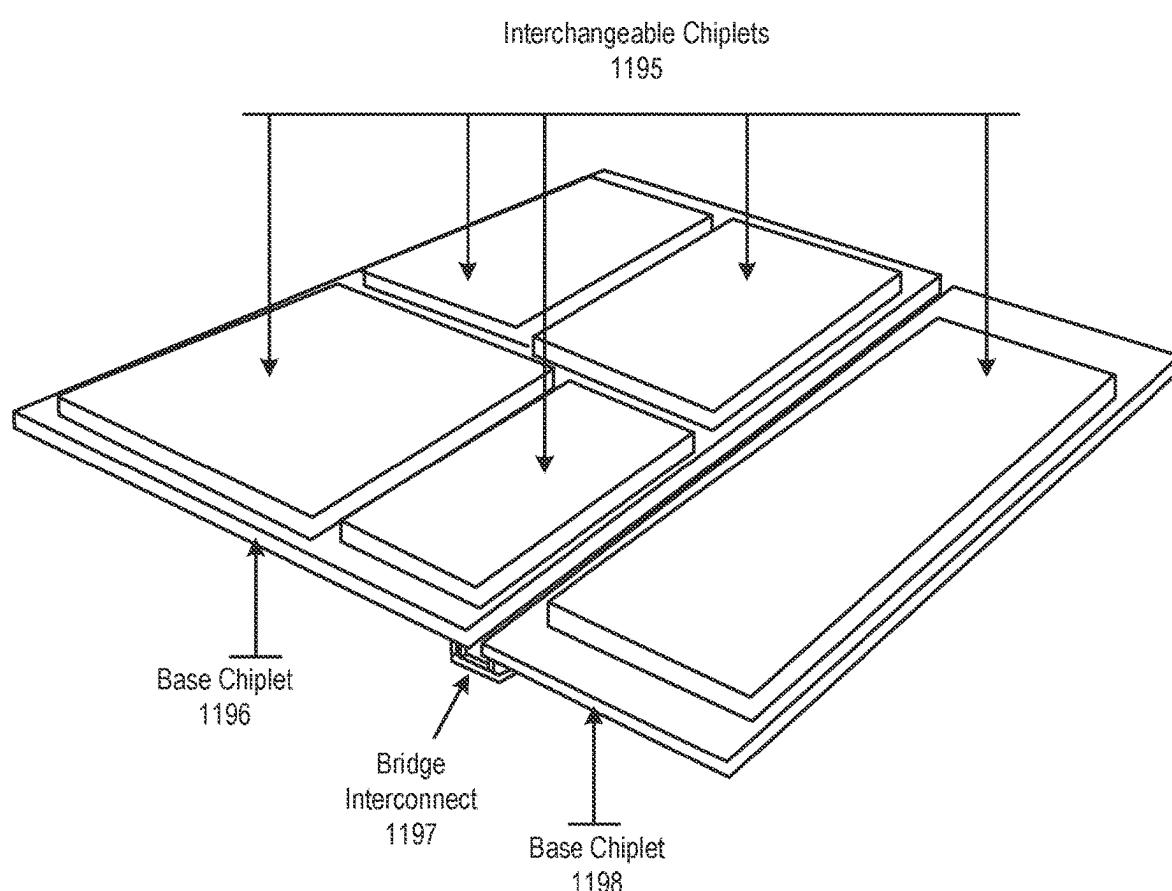


FIG. 11D

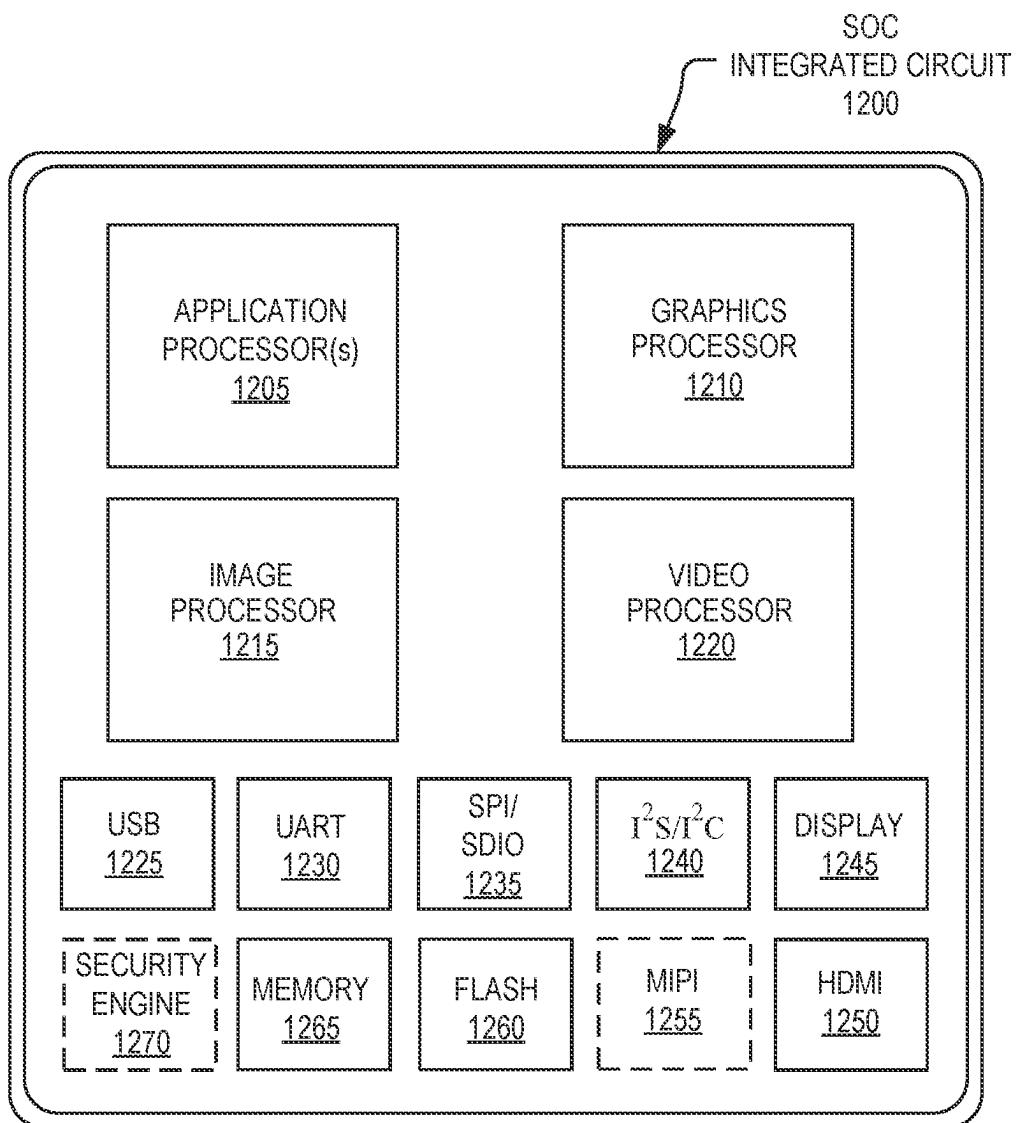


FIG. 12

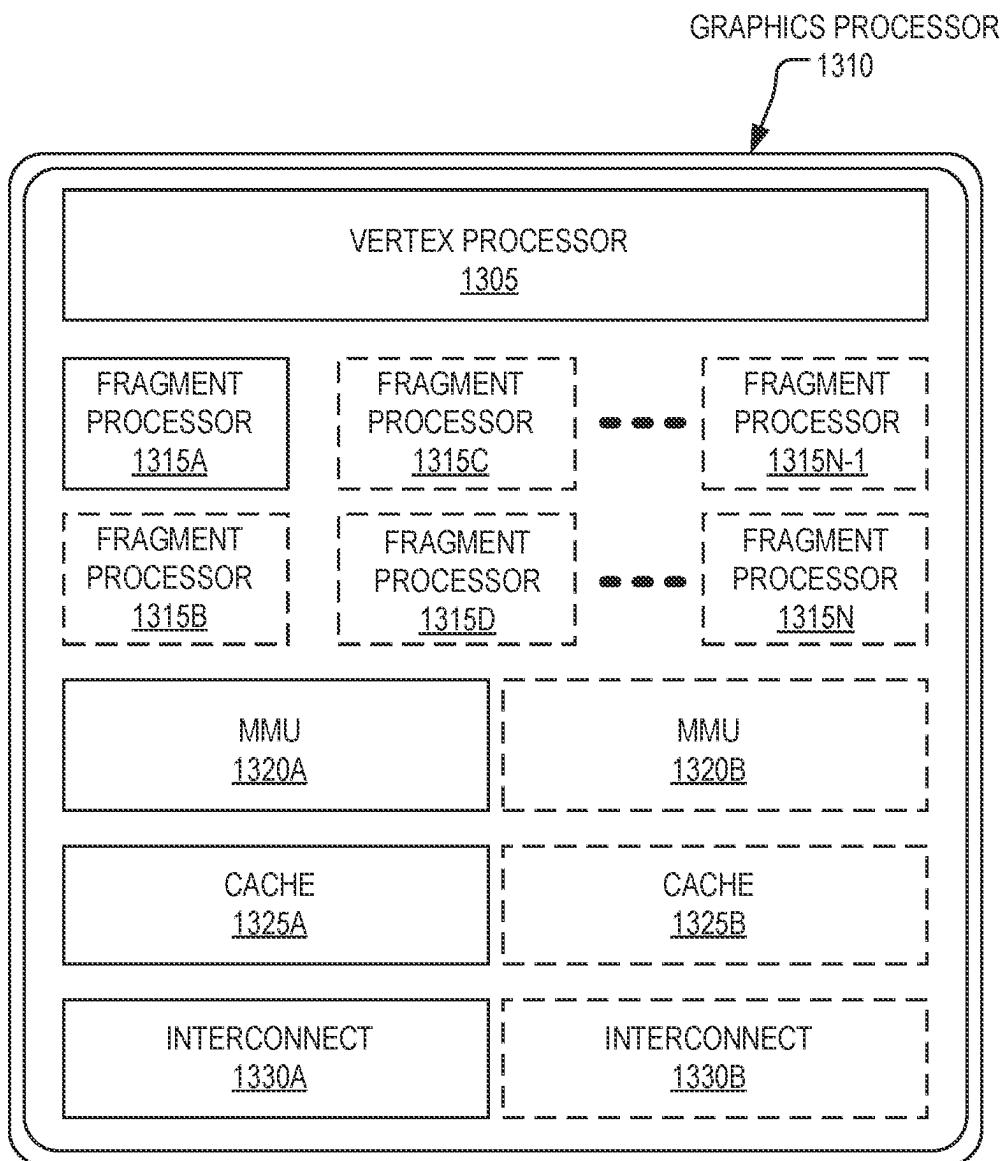


FIG. 13A

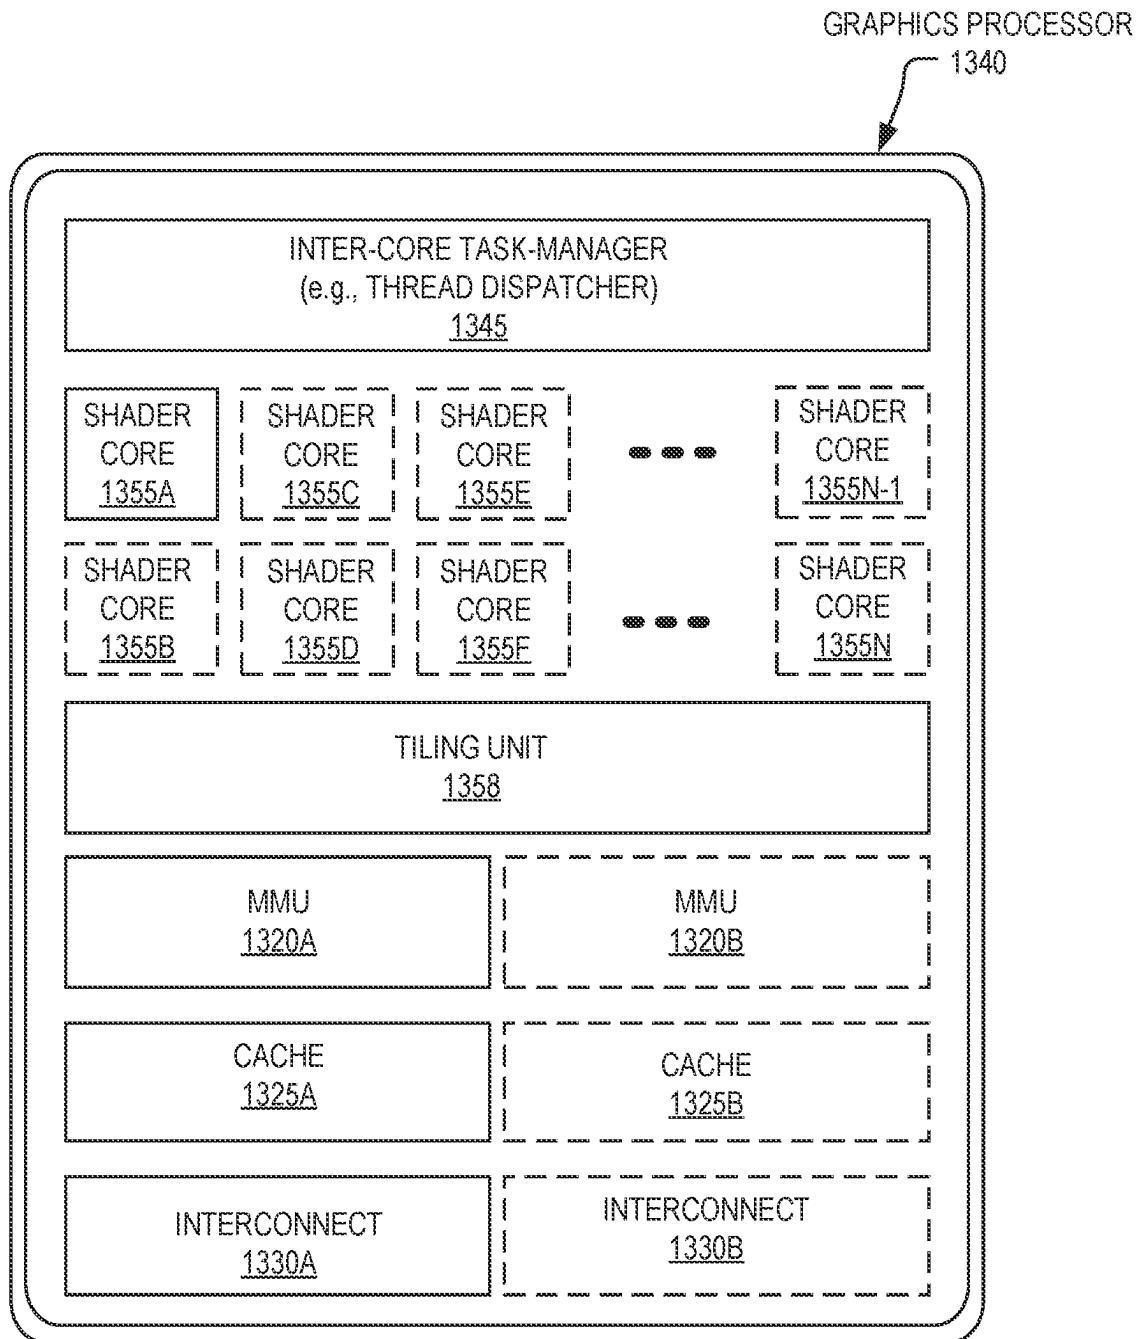


FIG. 13B

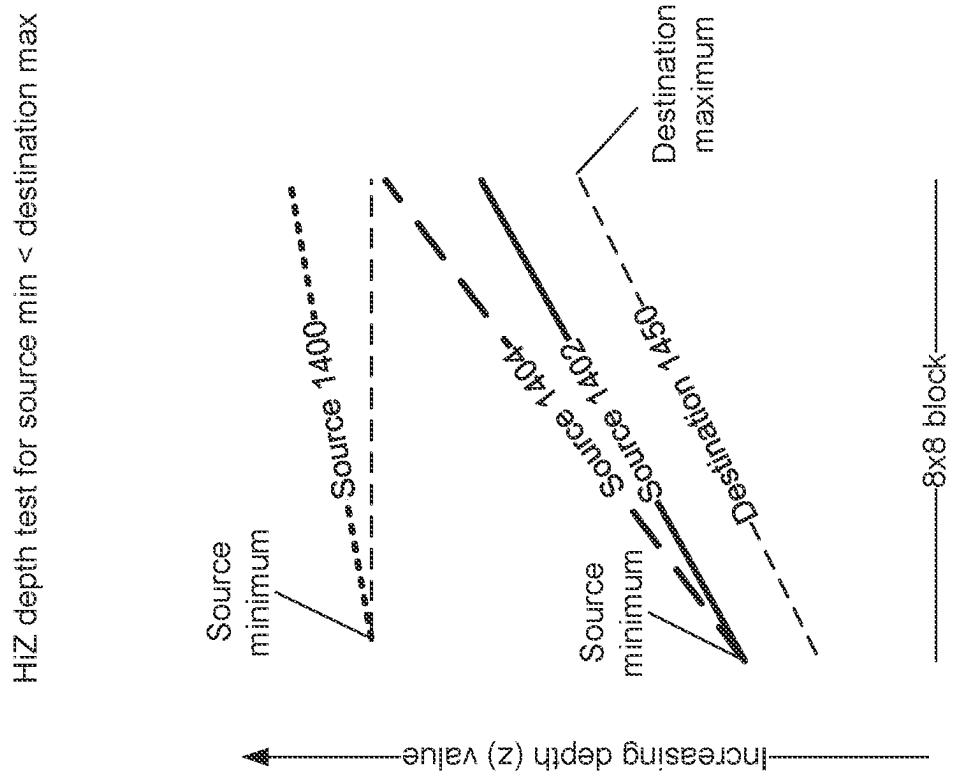


FIG. 14

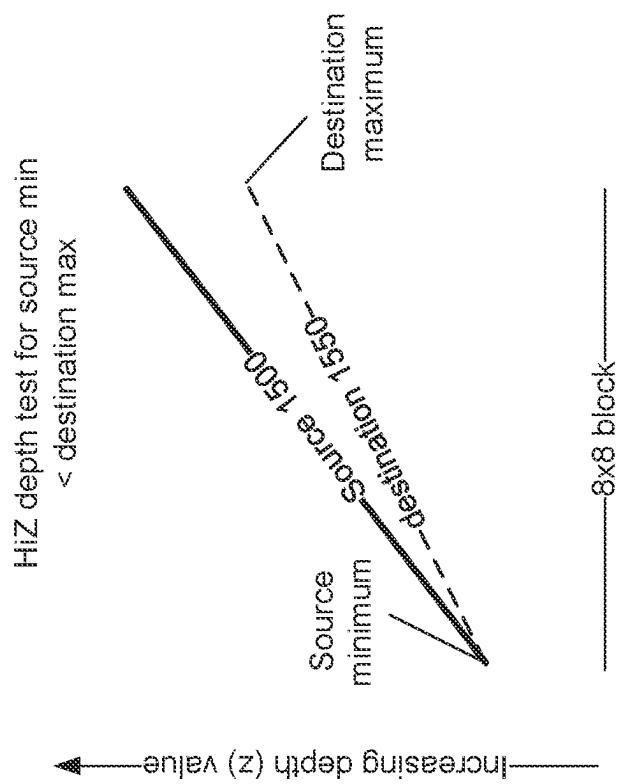


FIG. 15

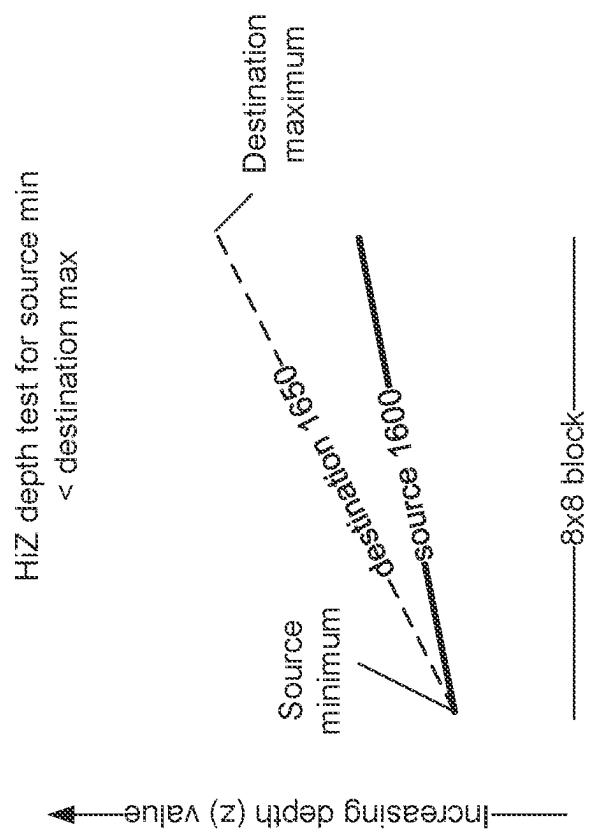


FIG. 16

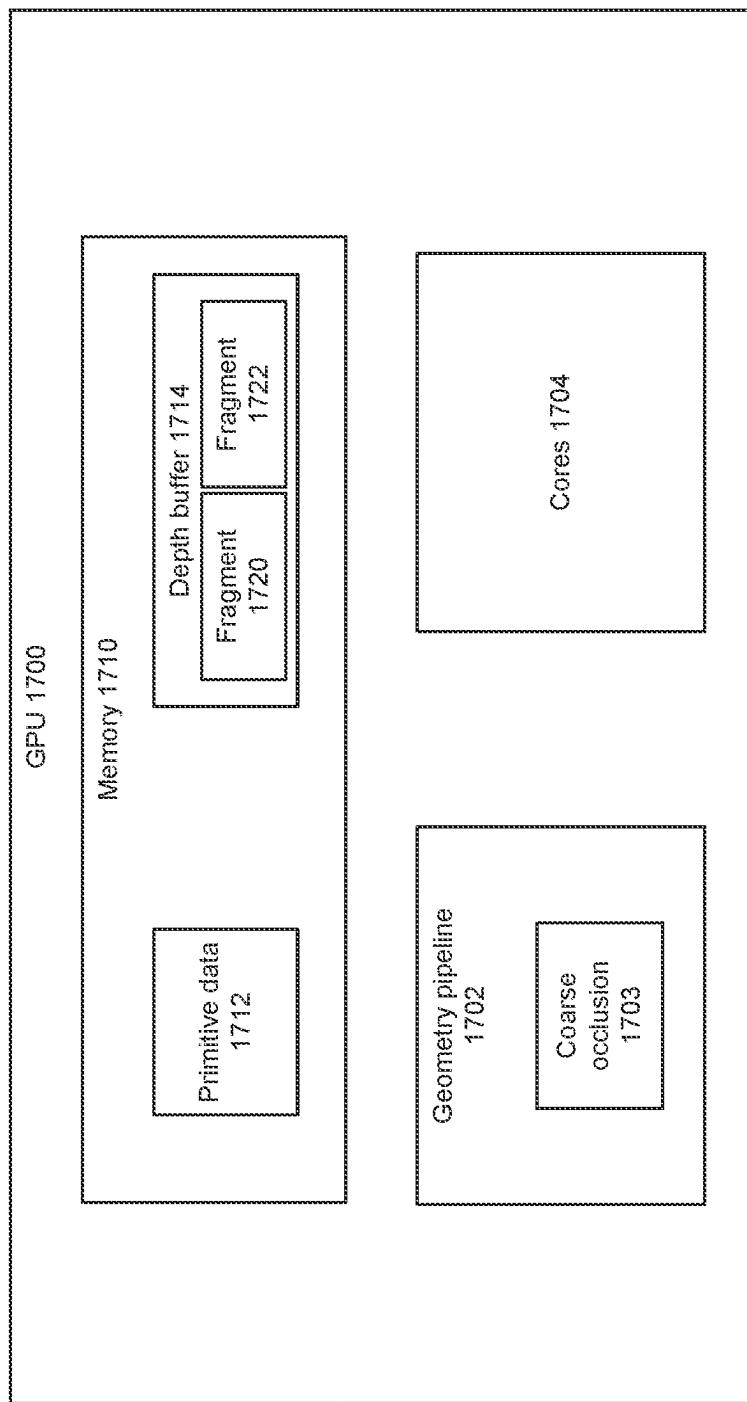
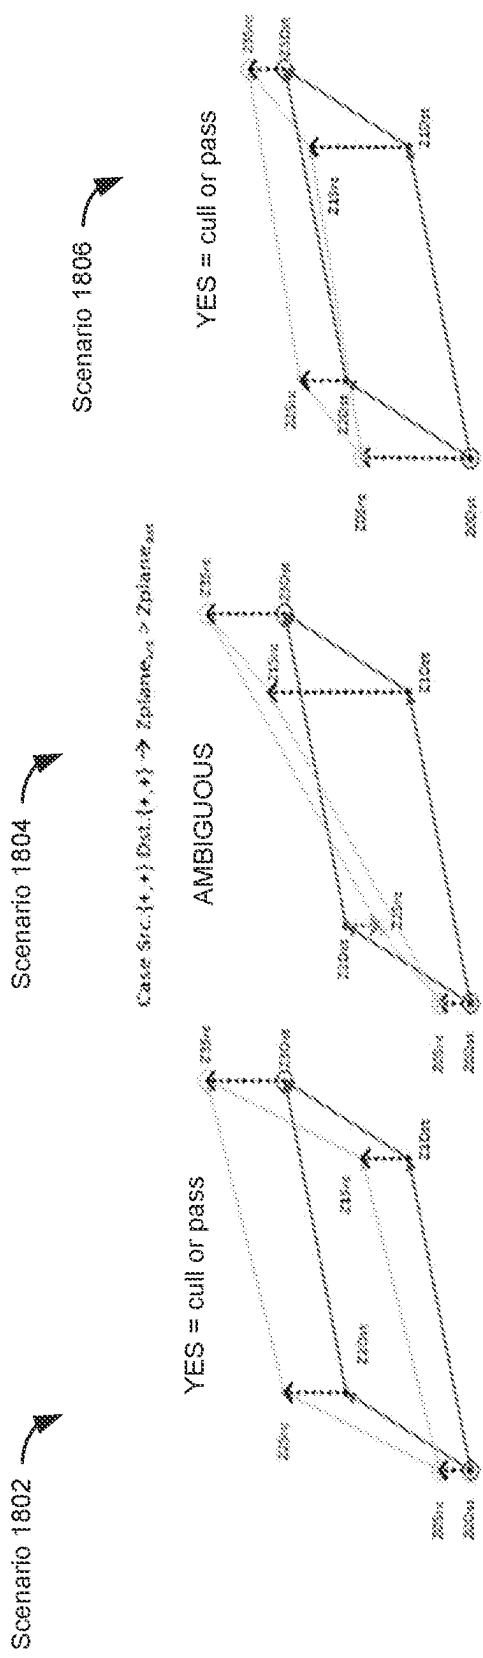
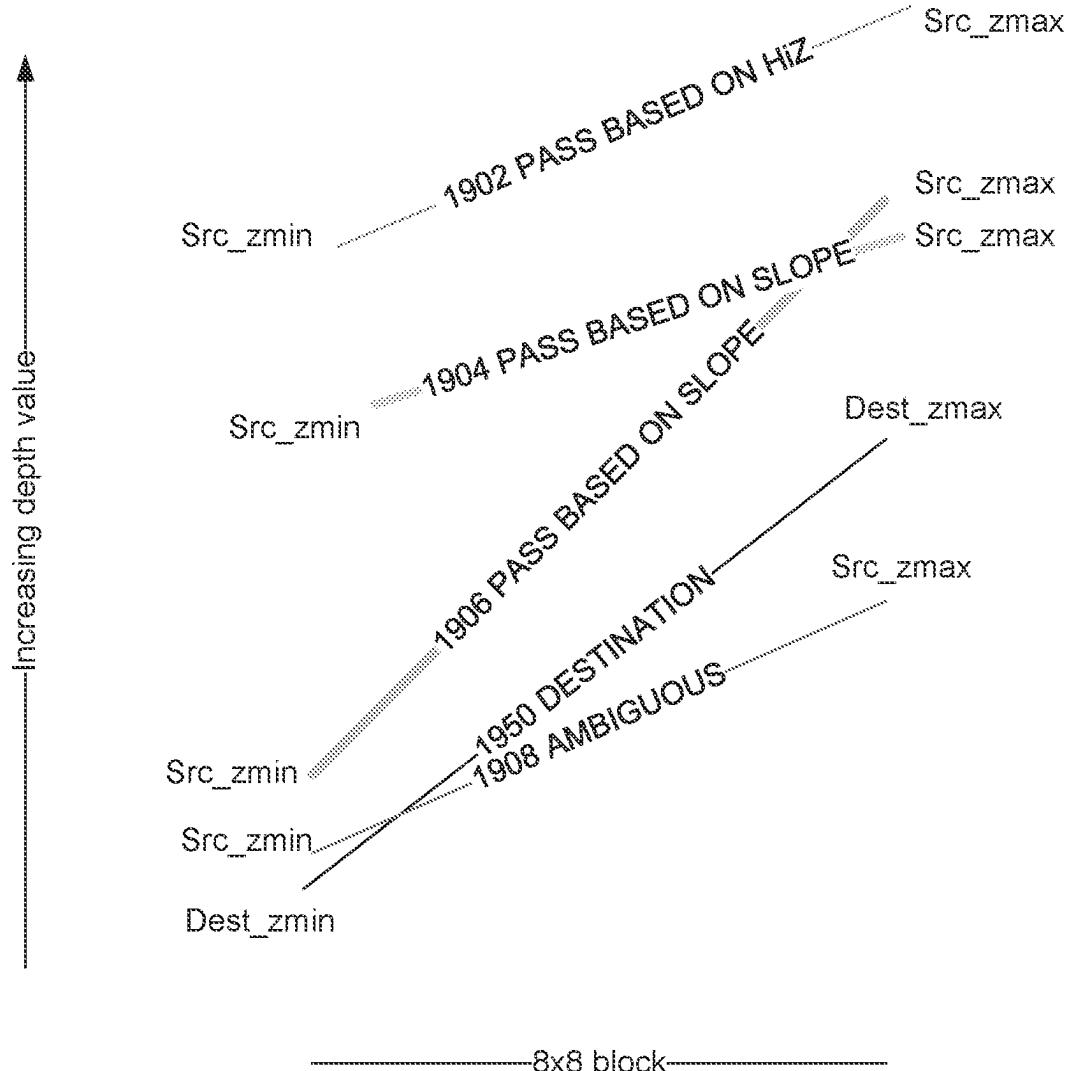


FIG. 17



१८



Depth test function = LESS
 $Z(x,y) = x^*Cx + y^*Cy + Cref$
Cx = slope in x dim
Cy = slope in y dim

FIG. 19

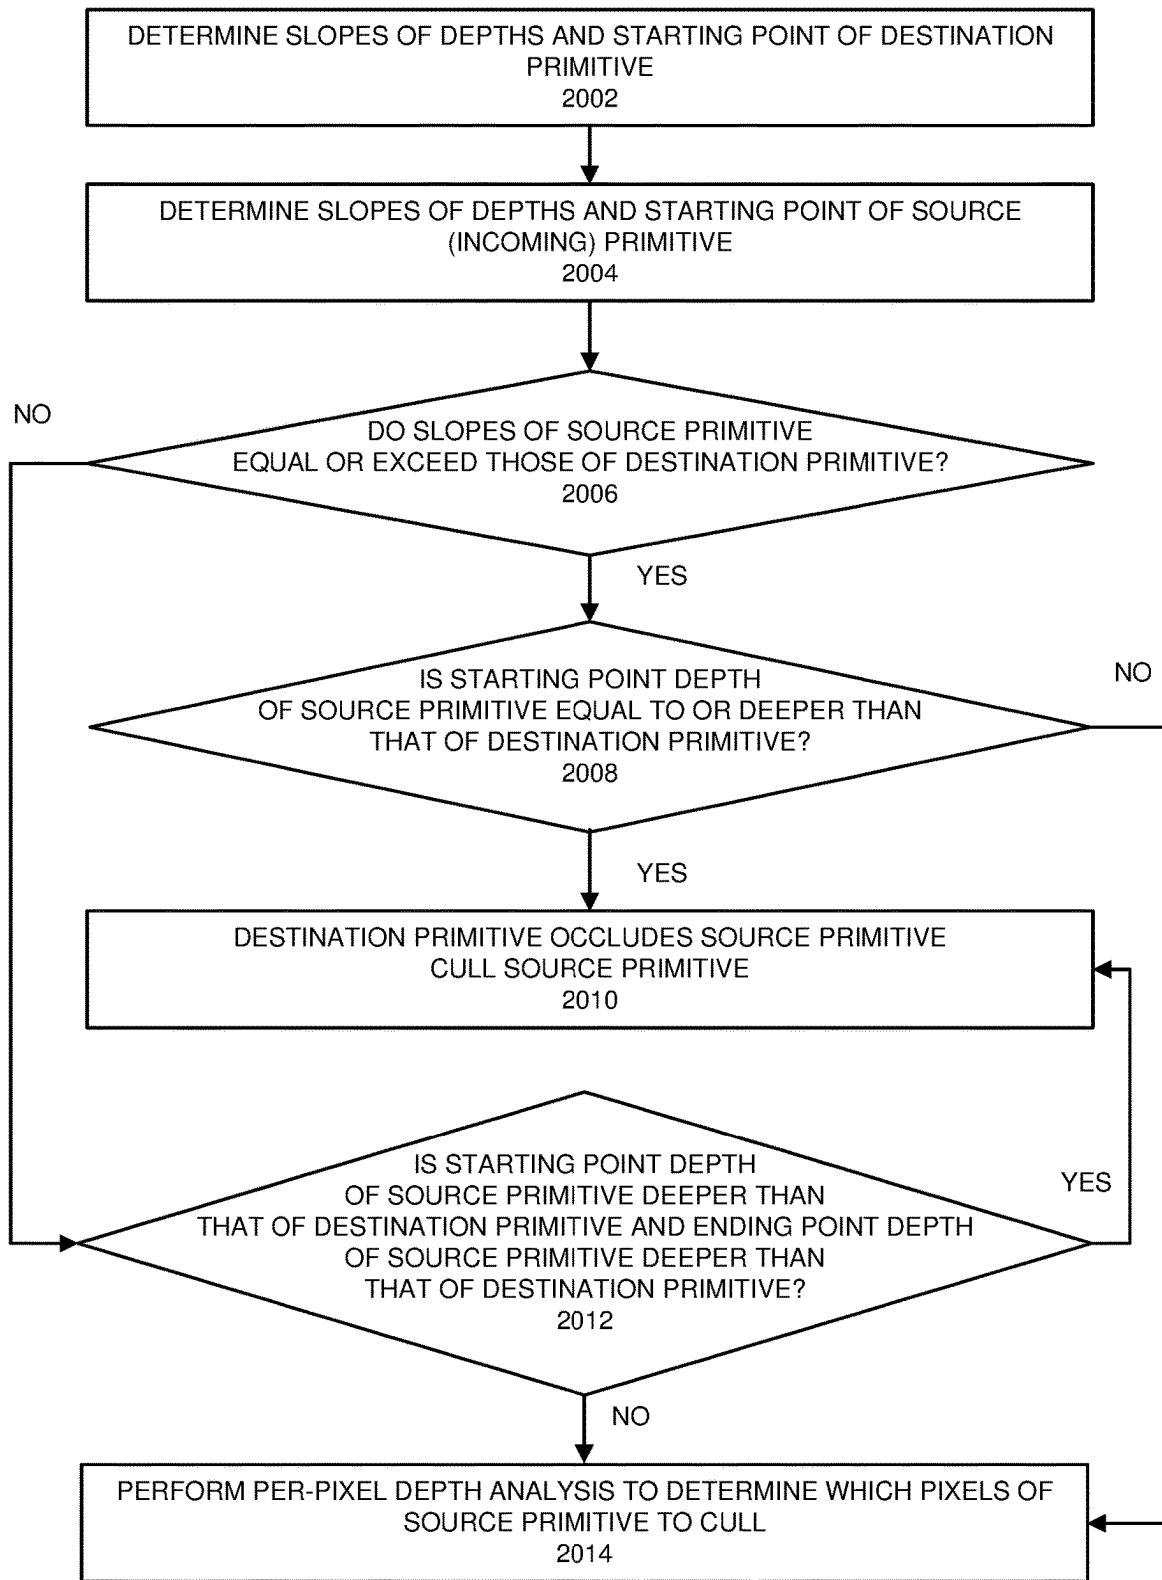


FIG. 20

OCCLUSION CULLING FOR DEPTH VALUES IN A DEPTH BUFFER

BACKGROUND

[0001] In the graphics processing unit (GPU) pipeline front-end, scene geometry is processed into primitives, such as triangles having associated vertex data. In the GPU pipeline back-end, depth and color values are assigned to pixels in screen space. Between the geometry and pixel pipelines is a rasterizer, which maps scene geometry to a pixel array, typically outputting tiles, fragments, or groups of a polygon. Polygons are ordered according to their distance from the viewer to properly project a three dimensional (3D) scene into two dimensional (2D) screen space. A 3D scene includes objects that are at least partially non-visible or occluded. For example, buildings and cars may occlude a number of objects in the scene.

[0002] A depth buffer or z-buffer is the last stage in a rendering pipeline where object fragments can be identified as occluded to avoid utilization of the GPU to render non-visible portions of object fragments. A depth buffer is an array built up from the depth values of one or more polygons. The depth (or z-value) of each polygon data group output by the rasterizer may be compared against a z-value stored in a depth buffer associated with the same location or position as the group. If the data group is closer to the viewpoint than the value stored in the depth buffer, it is rendered and the depth buffer is updated with the depth values associated with the newly generated data group. If the group is farther from the viewpoint than the value stored in the depth buffer, it is discarded as occluded. The depth buffer is therefore an array with each location in the array storing depth data of the polygon group that is visible from the viewpoint.

BRIEF DESCRIPTION OF THE DRAWINGS

[0003] So that the manner in which the features of the present embodiments can be understood in detail, a more particular description of the embodiments, briefly summarized above, may be had by reference to embodiments, some of which are illustrated in the appended drawings. It is to be noted, however, that the appended drawings illustrate only typical embodiments and are therefore not to be considered limiting of its scope.

[0004] FIG. 1 is a block diagram of a processing system, according to an embodiment.

[0005] FIG. 2A is a block diagram of an embodiment of a processor having one or more processor cores, an integrated memory controller, and an integrated graphics processor.

[0006] FIG. 2B is a block diagram of hardware logic of a graphics processor core block, according to some embodiments described herein.

[0007] FIG. 2C illustrates a graphics processing unit (GPU) that includes dedicated sets of graphics processing resources arranged into multi-core groups.

[0008] FIG. 2D is a block diagram of general-purpose graphics processing unit (GPGPU) that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein.

[0009] FIG. 3A is a block diagram of a graphics processor, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing

cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces.

[0010] FIG. 3B illustrates a graphics processor having a tiled architecture, according to embodiments described herein.

[0011] FIG. 3C illustrates a compute accelerator, according to embodiments described herein.

[0012] FIG. 4 is a block diagram of a graphics processing engine of a graphics processor in accordance with some embodiments.

[0013] FIG. 5A illustrates graphics core cluster, according to an embodiment.

[0014] FIG. 5B illustrates a vector engine of a graphics core, according to an embodiment.

[0015] FIG. 5C illustrates a matrix engine of a graphics core, according to an embodiment.

[0016] FIG. 6 illustrates a tile of a multi-tile processor, according to an embodiment.

[0017] FIG. 7 is a block diagram illustrating graphics processor instruction formats according to some embodiments.

[0018] FIG. 8 is a block diagram of another embodiment of a graphics processor.

[0019] FIG. 9A is a block diagram illustrating a graphics processor command format that may be used to program graphics processing pipelines according to some embodiments.

[0020] FIG. 9B is a block diagram illustrating a graphics processor command sequence according to an embodiment.

[0021] FIG. 10 illustrates an exemplary graphics software architecture for a data processing system according to some embodiments.

[0022] FIG. 11A is a block diagram illustrating an IP core development system that may be used to manufacture an integrated circuit to perform operations according to an embodiment.

[0023] FIG. 11B illustrates a cross-section side view of an integrated circuit package assembly, according to some embodiments described herein.

[0024] FIG. 11C illustrates a package assembly that includes multiple units of hardware logic chiplets connected to a substrate.

[0025] FIG. 11D illustrates a package assembly including interchangeable chiplets, according to an embodiment.

[0026] FIG. 12 is a block diagram illustrating an exemplary system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment.

[0027] FIG. 13A illustrates an exemplary graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment.

[0028] FIG. 13B illustrates an additional exemplary graphics processor of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment.

[0029] FIG. 14 depicts an example of a scenario in which a HiZ operation provides an ambiguous result for an occlusion analysis of a source primitive relative to a destination primitive.

[0030] FIG. 15 depicts an example of a scenario in which a HiZ operation provides an ambiguous result for an occlusion analysis of a source primitive relative to a destination primitive.

[0031] FIG. 16 depicts an example of a scenario in which a HiZ operation provides an ambiguous result for an occlusion analysis of source primitive relative to a destination primitive.

[0032] FIG. 17 depicts an example system.

[0033] FIG. 18 depicts an example of scenarios.

[0034] FIG. 19 depicts other examples of scenarios detected by slope-based occlusion.

[0035] FIG. 20 depicts an example process.

DETAILED DESCRIPTION

[0036] For occlusion analysis, in Microsoft® DirectX 11 and DirectX 12 pipelines, the GPU performs a coarse depth test before the per-pixel depth test. This coarse depth test can be based on a separate compressed depth buffer, which stores minimum/maximum ranges covering rectangular sections of the per pixel depth buffer, or a plane equation for entire chunk if a linear equation can denote depth at any given pixel for that chunk. For repeated overdraw to the depth buffer, where a new primitive is close in value to a previous plane equation, the HiZ test provides an ambiguous results, resulting in performing a computationally intensive and potentially slower per-pixel depth test, at full floating point precision, to determine whether the new primitive is either occluded or completely fills up relative to previous depth value at that location. This is an expensive operation constrained by the bandwidth and gates allocated to per pixel (or per-sample) depth test.

[0037] In some scenarios, a HiZ occlusion analysis can determine an ambiguous result that would utilize per-pixel depth testing to determine which pixels of a primitive are occluded. Various examples can determine whether an incoming primitive (source primitive) can be culled based on a slope of a source primitive compared to a slope of a second, destination, primitive. In some examples, if a slope of the source primitive in an X direction is greater than or equal to a slope of the destination primitive and the slope of the source primitive in the Y direction is greater than or equal to a slope of the destination primitive, the entire source primitive is occluded. In some examples, consideration of a starting reference depth of source and destination primitives in the X and Y directions can occur, in addition to a comparison of slopes in X and Y directions, to determine whether the entire source primitive is occluded.

System Overview

[0038] FIG. 1 is a block diagram of a processing system 100, according to an embodiment. Processing system 100 may be used in a single processor desktop system, a multiprocessor workstation system, or a server system having a large number of processors 102 or processor cores 107. In one embodiment, the processing system 100 is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices such as within Internet-of-things (IoT) devices with wired or wireless connectivity to a local or wide area network.

[0039] In one embodiment, processing system 100 can include, couple with, or be integrated within: a server-based gaming platform; a game console, including a game and media console; a mobile gaming console, a handheld game console, or an online game console. In some embodiments the processing system 100 is part of a mobile phone, smart

phone, tablet computing device or mobile Internet-connected device such as a laptop with low internal storage capacity. Processing system 100 can also include, couple with, or be integrated within: a wearable device, such as a smart watch wearable device; smart eyewear or clothing enhanced with augmented reality (AR) or virtual reality (VR) features to provide visual, audio or tactile outputs to supplement real world visual, audio or tactile experiences or otherwise provide text, audio, graphics, video, holographic images or video, or tactile feedback; other augmented reality (AR) device; or other virtual reality (VR) device. In some embodiments, the processing system 100 includes or is part of a television or set top box device. In one embodiment, processing system 100 can include, couple with, or be integrated within a self-driving vehicle such as a bus, tractor trailer, car, motor or electric power cycle, plane, or glider (or any combination thereof). The self-driving vehicle may use processing system 100 to process the environment sensed around the vehicle.

[0040] In some embodiments, the one or more processors 102 each include one or more processor cores 107 to process instructions which, when executed, perform operations for system or user software. In some embodiments, at least one of the one or more processor cores 107 is configured to process a specific instruction set 109. In some embodiments, instruction set 109 may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). One or more processor cores 107 may process a different instruction set 109, which may include instructions to facilitate the emulation of other instruction sets. Processor core 107 may also include other processing devices, such as a Digital Signal Processor (DSP).

[0041] In some embodiments, the processor 102 includes cache memory 104. Depending on the architecture, the processor 102 can have a single internal cache or multiple levels of internal cache. In some embodiments, the cache memory is shared among various components of the processor 102. In some embodiments, the processor 102 also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores 107 using known cache coherency techniques. A register file 106 can be additionally included in processor 102 and may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). Some registers may be general-purpose registers, while other registers may be specific to the design of the processor 102.

[0042] In some embodiments, one or more processor(s) 102 are coupled with one or more interface bus(es) 110 to transmit communication signals such as address, data, or control signals between processor 102 and other components in the processing system 100. The interface bus 110, in one embodiment, can be a processor bus, such as a version of the Direct Media Interface (DMI) bus. However, processor busses are not limited to the DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI express), memory busses, or other types of interface busses. In one embodiment the processor(s) 102 include a memory controller 116 and a platform controller hub 130. The memory controller 116 facilitates communication between a memory device and other components of the

processing system **100**, while the platform controller hub (PCH) **130** provides connections to I/O devices via a local I/O bus.

[0043] The memory device **120** can be a dynamic random-access memory (DRAM) device, a static random-access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In one embodiment the memory device **120** can operate as system memory for the processing system **100**, to store data **122** and instructions **121** for use when the one or more processors **102** executes an application or process. The memory controller **116** also couples with an optional external graphics processor **118**, which may communicate with the one or more graphics processors **108** in processors **102** to perform graphics and media operations. In some embodiments, graphics, media, and/or compute operations may be assisted by an accelerator **112** which is a coprocessor that can be configured to perform a specialized set of graphics, media, or compute operations. For example, in one embodiment the accelerator **112** is a matrix multiplication accelerator used to optimize machine learning or compute operations. In one embodiment the accelerator **112** is a ray-tracing accelerator that can be used to perform ray-tracing operations in concert with the graphics processor **108**. In one embodiment, an external accelerator **119** may be used in place of or in concert with the accelerator **112**.

[0044] In some embodiments a display device **111** can connect to the processor(s) **102**. The display device **111** can be one or more of an internal display device, as in a mobile electronic device or a laptop device or an external display device attached via a display interface (e.g., DisplayPort, etc.). In one embodiment the display device **111** can be a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

[0045] In some embodiments the platform controller hub **130** enables peripherals to connect to memory device **120** and processor **102** via a high-speed I/O bus. The I/O peripherals include, but are not limited to, an audio controller **146**, a network controller **134**, a firmware interface **128**, a wireless transceiver **126**, touch sensors **125**, a data storage device **124** (e.g., non-volatile memory, volatile memory, hard disk drive, flash memory, NAND, 3D NAND, 3D XPoint, etc.). The data storage device **124** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI express). The touch sensors **125** can include touch screen sensors, pressure sensors, or fingerprint sensors. The wireless transceiver **126** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, 5G, or Long-Term Evolution (LTE) transceiver. The firmware interface **128** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). The network controller **134** can enable a network connection to a wired network. In some embodiments, a high-performance network controller (not shown) couples with the interface bus **110**. The audio controller **146**, in one embodiment, is a multi-channel high-definition audio controller. In one embodiment the processing system **100** includes an optional legacy I/O controller **140** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to the system. The platform controller hub **130** can also connect to one or more Universal Serial Bus (USB)

controllers **142** connect input devices, such as keyboard and mouse **143** combinations, a camera **144**, or other USB input devices.

[0046] It will be appreciated that the processing system **100** shown is exemplary and not limiting, as other types of data processing systems that are differently configured may also be used. For example, an instance of the memory controller **116** and platform controller hub **130** may be integrated into a discrete external graphics processor, such as the external graphics processor **118**. In one embodiment the platform controller hub **130** and/or memory controller **116** may be external to the one or more processor(s) **102** and reside in a system chipset that is in communication with the processor(s) **102**.

[0047] For example, circuit boards ("sleds") can be used on which components such as CPUs, memory, and other components are placed and are designed for increased thermal performance. In some examples, processing components such as the processors are located on a top side of a sled while near memory, such as DIMMs, are located on a bottom side of the sled. As a result of the enhanced airflow provided by this design, the components may operate at higher frequencies and power levels than in typical systems, thereby increasing performance. Furthermore, the sleds are configured to blindly mate with power and data communication cables in a rack, thereby enhancing their ability to be quickly removed, upgraded, reinstalled, and/or replaced. Similarly, individual components located on the sleds, such as processors, accelerators, memory, and data storage drives, are configured to be easily upgraded due to their increased spacing from each other. In the illustrative embodiment, the components additionally include hardware attestation features to prove their authenticity.

[0048] A data center can utilize a single network architecture ("fabric") that supports multiple other network architectures including Ethernet and Omni-Path. The sleds can be coupled to switches via optical fibers, which provide higher bandwidth and lower latency than typical twisted pair cabling (e.g., Category 5, Category 5e, Category 6, etc.). Due to the high bandwidth, low latency interconnections and network architecture, the data center may, in use, pool resources, such as memory, accelerators (e.g., GPUs, graphics accelerators, FPGAs, ASICs, neural network and/or artificial intelligence accelerators, etc.), and data storage drives that are physically disaggregated, and provide them to compute resources (e.g., processors) on an as needed basis, enabling the compute resources to access the pooled resources as if they were local.

[0049] A power supply or source can provide voltage and/or current to processing system **100** or any component or system described herein. In one example, the power supply includes an AC to DC (alternating current to direct current) adapter to plug into a wall outlet. Such AC power can be renewable energy (e.g., solar power) power source. In one example, power source includes a DC power source, such as an external AC to DC converter. In one example, power source or power supply includes wireless charging hardware to charge via proximity to a charging field. In one example, power source can include an internal battery, alternating current supply, motion-based power supply, solar power supply, or fuel cell source.

[0050] FIGS. 2A-2D illustrate computing systems and graphics processors provided by embodiments described herein. The elements of FIGS. 2A-2D having the same

reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0051] FIG. 2A is a block diagram of an embodiment of a processor 200 having one or more processor cores 202A-202N, an integrated memory controller 214, and an integrated graphics processor 208. Processor 200 can include additional cores up to and including additional core 202N represented by the dashed lined boxes. Each of processor cores 202A-202N includes one or more internal cache units 204A-204N. In some embodiments each processor core also has access to one or more shared cached units 206. The internal cache units 204A-204N and shared cache units 206 represent a cache memory hierarchy within the processor 200. The cache memory hierarchy may include at least one level of instruction and data cache within each processor core and one or more levels of shared mid-level cache, such as a Level 2 (L2), Level 3 (L3), Level 4 (L4), or other levels of cache, where the highest level of cache before external memory is classified as the LLC. In some embodiments, cache coherency logic maintains coherency between the various cache units 206 and 204A-204N.

[0052] In some embodiments, processor 200 may also include a set of one or more bus controller units 216 and a system agent core 210. The one or more bus controller units 216 manage a set of peripheral buses, such as one or more PCI or PCI express busses. System agent core 210 provides management functionality for the various processor components. In some embodiments, system agent core 210 includes one or more integrated memory controllers 214 to manage access to various external memory devices (not shown).

[0053] In some embodiments, one or more of the processor cores 202A-202N include support for simultaneous multi-threading. In such embodiment, the system agent core 210 includes components for coordinating and operating cores 202A-202N during multi-threaded processing. System agent core 210 may additionally include a power control unit (PCU), which includes logic and components to regulate the power state of processor cores 202A-202N and graphics processor 208.

[0054] In some embodiments, processor 200 additionally includes graphics processor 208 to execute graphics processing operations. In some embodiments, the graphics processor 208 couples with the set of shared cache units 206, and the system agent core 210, including the one or more integrated memory controllers 214. In some embodiments, the system agent core 210 also includes a display controller 211 to drive graphics processor output to one or more coupled displays. In some embodiments, display controller 211 may also be a separate module coupled with the graphics processor via at least one interconnect, or may be integrated within the graphics processor 208.

[0055] In some embodiments, a ring-based interconnect 212 is used to couple the internal components of the processor 200. However, an alternative interconnect unit may be used, such as a point-to-point interconnect, a switched interconnect, a mesh interconnect, or other techniques, including techniques well known in the art. In some embodiments, graphics processor 208 couples with the ring-based interconnect 212 via an I/O link 213.

[0056] The exemplary I/O link 213 represents at least one of multiple varieties of I/O interconnects, including an on

package I/O interconnect which facilitates communication between various processor components and a high-performance embedded memory module 218, such as an eDRAM module or a high-bandwidth memory (HBM) module. In some embodiments, each of the processor cores 202A-202N and graphics processor 208 can use the embedded memory module 218 as a shared Last Level Cache.

[0057] In some embodiments, processor cores 202A-202N are homogenous cores executing the same instruction set architecture. In another embodiment, processor cores 202A-202N are heterogeneous in terms of instruction set architecture (ISA), where one or more of processor cores 202A-202N execute a first instruction set, while at least one of the other cores executes a subset of the first instruction set or a different instruction set. In one embodiment, processor cores 202A-202N are heterogeneous in terms of microarchitecture, where one or more cores having a relatively higher power consumption couple with one or more power cores having a lower power consumption. In one embodiment, processor cores 202A-202N are heterogeneous in terms of computational capability. Additionally, processor 200 can be implemented on one or more chips or as an SoC integrated circuit having the illustrated components, in addition to other components.

[0058] FIG. 2B is a block diagram of hardware logic of a graphics processor core block 219, according to some embodiments described herein. In some embodiments, elements of FIG. 2B having the same reference numbers (or names) as the elements of any other figure herein may operate or function in a manner similar to that described elsewhere herein. The graphics processor core block 219 is exemplary of one partition of a graphics processor. The graphics processor core block 219 can be included within the integrated graphics processor 208 of FIG. 2A or a discrete graphics processor, parallel processor, and/or compute accelerator. A graphics processor as described herein may include multiple graphics core blocks based on target power and performance envelopes. Each graphics processor core block 219 can include a function block 230 coupled with multiple graphics cores 221A-221F that include modular blocks of fixed function logic and general-purpose programmable logic. The graphics processor core block 219 also includes shared/cache memory 236 that is accessible by all graphics cores 221A-221F, rasterizer logic 237, and additional fixed function logic 238.

[0059] In some embodiments, the function block 230 includes a geometry/fixed function pipeline 231 that can be shared by all graphics cores in the graphics processor core block 219. In various embodiments, the geometry/fixed function pipeline 231 includes a 3D geometry pipeline a video front-end unit, a thread spawner and global thread dispatcher, and a unified return buffer manager, which manages unified return buffers. In one embodiment the function block 230 also includes a graphics SoC interface 232, a graphics microcontroller 233, and a media pipeline 234. The graphics SoC interface 232 provides an interface between the graphics processor core block 219 and other core blocks within a graphics processor or compute accelerator SoC. The graphics microcontroller 233 is a programmable sub-processor that is configurable to manage various functions of the graphics processor core block 219, including thread dispatch, scheduling, and pre-emption. The media pipeline 234 includes logic to facilitate the decoding, encoding, pre-processing, and/or post-processing of multimedia

data, including image and video data. The media pipeline 234 implement media operations via requests to compute or sampling logic within the graphics cores 221-221F. One or more pixel backends 235 can also be included within the function block 230. The pixel backends 235 include a cache memory to store pixel color values and can perform blend operations and lossless color compression of rendered pixel data.

[0060] In one embodiment the graphics SoC interface 232 enables the graphics processor core block 219 to communicate with general-purpose application processor cores (e.g., CPUs) and/or other components within an SoC or a system host CPU that is coupled with the SoC via a peripheral interface. The graphics SoC interface 232 also enables communication with off-chip memory hierarchy elements such as a shared last level cache memory, system RAM, and/or embedded on-chip or on-package DRAM. The SoC interface 232 can also enable communication with fixed function devices within the SoC, such as camera imaging pipelines, and enables the use of and/or implements global memory atomics that may be shared between the graphics processor core block 219 and CPUs within the SoC. The graphics SoC interface 232 can also implement power management controls for the graphics processor core block 219 and enable an interface between a clock domain of the graphics processor core block 219 and other clock domains within the SoC. In one embodiment the graphics SoC interface 232 enables receipt of command buffers from a command streamer and global thread dispatcher that are configured to provide commands and instructions to each of one or more graphics cores within a graphics processor. The commands and instructions can be dispatched to the media pipeline 234 when media operations are to be performed, the geometry and fixed function pipeline 231 when graphics processing operations are to be performed. When compute operations are to be performed, compute dispatch logic can dispatch the commands to the graphics cores 221A-221F, bypassing the geometry and media pipelines.

[0061] The graphics microcontroller 233 can be configured to perform various scheduling and management tasks for the graphics processor core block 219. In one embodiment the graphics microcontroller 233 can perform graphics and/or compute workload scheduling on the various vector engines 222A-222F, 224A-224F and matrix engines 223A-223F, 225A-225F within the graphics cores 221A-221F. In this scheduling model, host software executing on a CPU core of an SoC including the graphics processor core block 219 can submit workloads to one of multiple graphic processor doorbells, which invokes a scheduling operation on the appropriate graphics engine. Scheduling operations include determining which workload to run next, submitting a workload to a command streamer, pre-empting existing workloads running on an engine, monitoring progress of a workload, and notifying host software when a workload is complete. In one embodiment the graphics microcontroller 233 can also facilitate low-power or idle states for the graphics processor core block 219, providing the graphics processor core block 219 with the ability to save and restore registers within the graphics processor core block 219 across low-power state transitions independently from the operating system and/or graphics driver software on the system.

[0062] The graphics processor core block 219 may have greater than or fewer than the illustrated graphics cores 221A-221F, up to N modular graphics cores. For each set of

N graphics cores, the graphics processor core block 219 can also include shared/cache memory 236, which can be configured as shared memory or cache memory, rasterizer logic 237, and additional fixed function logic 238 to accelerate various graphics and compute processing operations.

[0063] Within each graphics cores 221A-221F is set of execution resources that may be used to perform graphics, media, and compute operations in response to requests by graphics pipeline, media pipeline, or shader programs. The graphics cores 221A-221F include multiple vector engines 222A-222F, 224A-224F, matrix acceleration units 223A-223F, 225A-225D, cache/shared local memory (SLM), a sampler 226A-226F, and a ray tracing unit 227A-227F.

[0064] The vector engines 222A-222F, 224A-224F are general-purpose graphics processing units capable of performing floating-point and integer/fixed-point logic operations in service of a graphics, media, or compute operation, including graphics, media, or compute/GPGPU programs. The vector engines 222A-222F, 224A-224F can operate at variable vector widths using SIMD, SIMT, or SIMT+SIMD execution modes. The matrix acceleration units 223A-223F, 225A-225D include matrix-matrix and matrix-vector acceleration logic that improves performance on matrix operations, particularly low and mixed precision (e.g., INT8, FP16, BF16) matrix operations used for machine learning. In one embodiment, each of the matrix acceleration units 223A-223F, 225A-225D includes one or more systolic arrays of processing elements that can perform concurrent matrix multiply or dot product operations on matrix elements.

[0065] The sampler 226A-226F can read media or texture data into memory and can sample data differently based on a configured sampler state and the texture/media format that is being read. Threads executing on the vector engines 222A-222F, 224A-224F or matrix acceleration units 223A-223F, 225A-225D can make use of the cache/SLM 228A-228F within each execution core. The cache/SLM 228A-228F can be configured as cache memory or as a pool of shared memory that is local to each of the respective graphics cores 221A-221F. The ray tracing units 227A-227F within the graphics cores 221A-221F include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. In one embodiment the ray tracing units 227A-227F include circuitry for performing depth testing and culling (e.g., using a depth buffer or similar arrangement). In one implementation, the ray tracing units 227A-227F perform traversal and intersection operations in concert with image denoising, at least a portion of which may be performed using an associated matrix acceleration unit 223A-223F, 225A-225D.

[0066] FIG. 2C illustrates a graphics processing unit (GPU) 239 that includes dedicated sets of graphics processing resources arranged into multi-core groups 240A-240N. The details of multi-core group 240A are illustrated. Multi-core groups 240B-240N may be equipped with the same or similar sets of graphics processing resources.

[0067] As illustrated, a multi-core group 240A may include a set of graphics cores 243, a set of tensor cores 244, and a set of ray tracing cores 245. A scheduler/dispatcher 241 schedules and dispatches the graphics threads for execution on the various cores 243, 244, 245. In one embodiment the tensor cores 244 are sparse tensor cores with hardware

to enable multiplication operations having a zero-value input to be bypassed. The graphics cores 243 of the GPU 239 of FIG. 2C differ in hierarchical abstraction level relative to the graphics cores 221A-221F of FIG. 2B, which are analogous to the multi-core groups 240A-240N of FIG. 2C. The graphics cores 243, tensor cores 244, and ray tracing cores 245 of FIG. 2C are analogous to, respectively, the vector engines 222A-222F, 224A-224F, matrix engines 223A-223F, 225A-225F, and ray tracing units 227A-227F of FIG. 2B.

[0068] A set of register files 242 can store operand values used by the cores 243, 244, 245 when executing the graphics threads. These may include, for example, integer registers for storing integer values, floating point registers for storing floating point values, vector registers for storing packed data elements (integer and/or floating-point data elements) and tile registers for storing tensor/matrix values. In one embodiment, the tile registers are implemented as combined sets of vector registers.

[0069] One or more combined level 1 (L1) caches and shared memory units 247 store graphics data such as texture data, vertex data, pixel data, ray data, bounding volume data, etc., locally within each multi-core group 240A. One or more texture units 247 can also be used to perform texturing operations, such as texture mapping and sampling. A Level 2 (L2) cache 253 shared by all or a subset of the multi-core groups 240A-240N stores graphics data and/or instructions for multiple concurrent graphics threads. As illustrated, the L2 cache 253 may be shared across a plurality of multi-core groups 240A-240N. One or more memory controllers 248 couple the GPU 239 to a memory 249 which may be a system memory (e.g., DRAM) and/or a dedicated graphics memory (e.g., GDDR6 memory).

[0070] Input/output (I/O) circuitry 250 couples the GPU 239 to one or more I/O devices 252 such as digital signal processors (DSPs), network controllers, or user input devices. An on-chip interconnect may be used to couple the I/O devices 252 to the GPU 239 and memory 249. One or more I/O memory management units (IOMMUs) 251 of the I/O circuitry 250 couple the I/O devices 252 directly to the memory 249. In one embodiment, the IOMMU 251 manages multiple sets of page tables to map virtual addresses to physical addresses in memory 249. In this embodiment, the I/O devices 252, CPU(s) 246, and GPU 239 may share the same virtual address space.

[0071] In one implementation, the IOMMU 251 supports virtualization. In this case, it may manage a first set of page tables to map guest/graphics virtual addresses to guest/graphics physical addresses and a second set of page tables to map the guest/graphics physical addresses to system/host physical addresses (e.g., within memory 249). The base addresses of each of the first and second sets of page tables may be stored in control registers and swapped out on a context switch (e.g., so that the new context is provided with access to the relevant set of page tables). While not illustrated in FIG. 2C, each of the cores 243, 244, 245 and/or multi-core groups 240A-240N may include translation lookaside buffers (TLBs) to cache guest virtual to guest physical translations, guest physical to host physical translations, and guest virtual to host physical translations.

[0072] In one embodiment, the CPUs 246, GPU 239, and I/O devices 252 are integrated on a single semiconductor chip and/or chip package. The memory 249 may be integrated on the same chip or may be coupled to the memory

controllers 248 via an off-chip interface. In one implementation, the memory 249 comprises GDDR6 memory which shares the same virtual address space as other physical system-level memories, although the underlying principles of the embodiments described herein are not limited to this specific implementation.

[0073] In one embodiment, the tensor cores 244 include a plurality of functional units specifically designed to perform matrix operations, which are the fundamental compute operation used to perform deep learning operations. For example, simultaneous matrix multiplication operations may be used for neural network training and inferencing. The tensor cores 244 may perform matrix processing using a variety of operand precisions including single precision floating-point (e.g., 32 bits), half-precision floating point (e.g., 16 bits), integer words (16 bits), bytes (8 bits), and half-bytes (4 bits). In one embodiment, a neural network implementation extracts features of each rendered scene, potentially combining details from multiple frames, to construct a high-quality final image.

[0074] In deep learning implementations, parallel matrix multiplication work may be scheduled for execution on the tensor cores 244. The training of neural networks, in particular, requires a significant number of matrix dot product operations. In order to process an inner-product formulation of an NxNxN matrix multiply, the tensor cores 244 may include at least N dot-product processing elements. Before the matrix multiply begins, one entire matrix is loaded into tile registers and at least one column of a second matrix is loaded each cycle for N cycles. Each cycle, there are N dot products that are processed.

[0075] Matrix elements may be stored at different precisions depending on the particular implementation, including 16-bit words, 8-bit bytes (e.g., INT8) and 4-bit half-bytes (e.g., INT4). Different precision modes may be specified for the tensor cores 244 to ensure that the most efficient precision is used for different workloads (e.g., such as inferencing workloads which can tolerate quantization to bytes and half-bytes).

[0076] In one embodiment, the ray tracing cores 245 accelerate ray tracing operations for both real-time ray tracing and non-real-time ray tracing implementations. In particular, the ray tracing cores 245 include ray traversal/intersection circuitry for performing ray traversal using bounding volume hierarchies (BVHs) and identifying intersections between rays and primitives enclosed within the BVH volumes. The ray tracing cores 245 may also include circuitry for performing depth testing and culling (e.g., using a Z buffer or similar arrangement). In one implementation, the ray tracing cores 245 perform traversal and intersection operations in concert with the image denoising techniques described herein, at least a portion of which may be executed on the tensor cores 244. For example, in one embodiment, the tensor cores 244 implement a deep learning neural network to perform denoising of frames generated by the ray tracing cores 245. However, the CPU(s) 246, graphics cores 243, and/or ray tracing cores 245 may also implement all or a portion of the denoising and/or deep learning algorithms.

[0077] In addition, as described above, a distributed approach to denoising may be employed in which the GPU 239 is in a computing device coupled to other computing devices over a network or high-speed interconnect. In this embodiment, the interconnected computing devices share neural network learning/training data to improve the speed

with which the overall system learns to perform denoising for different types of image frames and/or different graphics applications.

[0078] In one embodiment, the ray tracing cores 245 process all BVH traversal and ray-primitive intersections, saving the graphics cores 243 from being overloaded with thousands of instructions per ray. In one embodiment, each ray tracing core 245 includes a first set of specialized circuitry for performing bounding box tests (e.g., for traversal operations) and a second set of specialized circuitry for performing the ray-triangle intersection tests (e.g., intersecting rays which have been traversed). Thus, in one embodiment, the multi-core group 240A can simply launch a ray probe, and the ray tracing cores 245 independently perform ray traversal and intersection and return hit data (e.g., a hit, no hit, multiple hits, etc.) to the thread context. The other cores 243, 244 are freed to perform other graphics or compute work while the ray tracing cores 245 perform the traversal and intersection operations.

[0079] In one embodiment, each ray tracing core 245 includes a traversal unit to perform BVH testing operations and an intersection unit which performs ray-primitive intersection tests. The intersection unit generates a “hit”, “no hit”, or “multiple hit” response, which it provides to the appropriate thread. During the traversal and intersection operations, the execution resources of the other cores (e.g., graphics cores 243 and tensor cores 244) are freed to perform other forms of graphics work.

[0080] In one particular embodiment described below, a hybrid rasterization/ray tracing approach is used in which work is distributed between the graphics cores 243 and ray tracing cores 245.

[0081] In one embodiment, the ray tracing cores 245 (and/or other cores 243, 244) include hardware support for a ray tracing instruction set such as Microsoft’s DirectX Ray Tracing (DXR) which includes a DispatchRays command, as well as ray-generation, closest-hit, any-hit, and miss shaders, which enable the assignment of unique sets of shaders and textures for each object. Another ray tracing platform which may be supported by the ray tracing cores 245, graphics cores 243 and tensor cores 244 is Vulkan 1.1.85. Note, however, that the underlying principles of the embodiments described herein are not limited to any particular ray tracing ISA.

[0082] In general, the various cores 245, 244, 243 may support a ray tracing instruction set that includes instructions/functions for ray generation, closest hit, any hit, ray-primitive intersection, per-primitive and hierarchical bounding box construction, miss, visit, and exceptions. More specifically, one embodiment includes ray tracing instructions to perform the following functions:

[0083] Ray Generation—Ray generation instructions may be executed for each pixel, sample, or other user-defined work assignment.

[0084] Closest Hit—A closest hit instruction may be executed to locate the closest intersection point of a ray with primitives within a scene.

[0085] Any Hit—An any hit instruction identifies multiple intersections between a ray and primitives within a scene, potentially to identify a new closest intersection point.

[0086] Intersection—An intersection instruction performs a ray-primitive intersection test and outputs a result.

[0087] Per-primitive Bounding box Construction—This instruction builds a bounding box around a given primitive

or group of primitives (e.g., when building a new BVH or other acceleration data structure).

[0088] Miss—Indicates that a ray misses all geometry within a scene, or specified region of a scene.

[0089] Visit—Indicates the child volumes a ray will traverse.

[0090] Exceptions—Includes various types of exception handlers (e.g., invoked for various error conditions).

[0091] In one embodiment the ray tracing cores 245 may be adapted to accelerate general-purpose compute operations that can be accelerated using computational techniques that are analogous to ray intersection tests. A compute framework can be provided that enables shader programs to be compiled into low level instructions and/or primitives that perform general-purpose compute operations via the ray tracing cores. Exemplary computational problems that can benefit from compute operations performed on the ray tracing cores 245 include computations involving beam, wave, ray, or particle propagation within a coordinate space. Interactions associated with that propagation can be computed relative to a geometry or mesh within the coordinate space. For example, computations associated with electromagnetic signal propagation through an environment can be accelerated via the use of instructions or primitives that are executed via the ray tracing cores. Diffraction and reflection of the signals by objects in the environment can be computed as direct ray-tracing analogies.

[0092] Ray tracing cores 245 can also be used to perform computations that are not directly analogous to ray tracing. For example, mesh projection, mesh refinement, and volume sampling computations can be accelerated using the ray tracing cores 245. Generic coordinate space calculations, such as nearest neighbor calculations can also be performed. For example, the set of points near a given point can be discovered by defining a bounding box in the coordinate space around the point. BVH and ray probe logic within the ray tracing cores 245 can then be used to determine the set of point intersections within the bounding box. The intersections constitute the origin point and the nearest neighbors to that origin point. Computations that are performed using the ray tracing cores 245 can be performed in parallel with computations performed on the graphics cores 243 and tensor cores 244. A shader compiler can be configured to compile a compute shader or other general-purpose graphics processing program into low level primitives that can be parallelized across the graphics cores 243, tensor cores 244, and ray tracing cores 245.

[0093] FIG. 2D is a block diagram of general-purpose graphics processing unit (GPGPU) 270 that can be configured as a graphics processor and/or compute accelerator, according to embodiments described herein. The GPGPU 270 can interconnect with host processors (e.g., one or more CPU(s) 246) and memory 271, 272 via one or more system and/or memory busses. In one embodiment the memory 271 is system memory that may be shared with the one or more CPU(s) 246, while memory 272 is device memory that is dedicated to the GPGPU 270. In one embodiment, components within the GPGPU 270 and memory 272 may be mapped into memory addresses that are accessible to the one or more CPU(s) 246. Access to memory 271 and 272 may be facilitated via a memory controller 268. In one embodiment the memory controller 268 includes an internal direct

memory access (DMA) controller 269 or can include logic to perform operations that would otherwise be performed by a DMA controller.

[0094] The GPGPU 270 includes multiple cache memories, including an L2 cache 253, L1 cache 254, an instruction cache 255, and shared memory 256, at least a portion of which may also be partitioned as a cache memory. The GPGPU 270 also includes multiple compute units 260A-260N, which represent a hierarchical abstraction level analogous to the graphics cores 221A-221F of FIG. 2B and the multi-core groups 240A-240N of FIG. 2C. Each compute unit 260A-260N includes a set of vector registers 261, scalar registers 262, vector logic units 263, and scalar logic units 264. The compute units 260A-260N can also include local shared memory 265 and a program counter 266. The compute units 260A-260N can couple with a constant cache 267, which can be used to store constant data, which is data that will not change during the run of kernel or shader program that executes on the GPGPU 270. In one embodiment the constant cache 267 is a scalar data cache and cached data can be fetched directly into the scalar registers 262.

[0095] During operation, the one or more CPU(s) 246 can write commands into registers or memory in the GPGPU 270 that has been mapped into an accessible address space. The command processors 257 can read the commands from registers or memory and determine how those commands will be processed within the GPGPU 270. A thread dispatcher 258 can then be used to dispatch threads to the compute units 260A-260N to perform those commands. Each compute unit 260A-260N can execute threads independently of the other compute units. Additionally, each compute unit 260A-260N can be independently configured for conditional computation and can conditionally output the results of computation to memory. The command processors 257 can interrupt the one or more CPU(s) 246 when the submitted commands are complete.

[0096] FIGS. 3A-3C illustrate block diagrams of additional graphics processor and compute accelerator architectures provided by embodiments described herein. The elements of FIGS. 3A-3C having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0097] FIG. 3A is a block diagram of a graphics processor 300, which may be a discrete graphics processing unit, or may be a graphics processor integrated with a plurality of processing cores, or other semiconductor devices such as, but not limited to, memory devices or network interfaces. In some embodiments, the graphics processor communicates via a memory mapped I/O interface to registers on the graphics processor and with commands placed into the processor memory. In some embodiments, graphics processor 300 includes a memory interface 314 to access memory. Memory interface 314 can be an interface to local memory, one or more internal caches, one or more shared external caches, and/or to system memory.

[0098] In some embodiments, graphics processor 300 also includes a display controller 302 to drive display output data to a display device 318. Display controller 302 includes hardware for one or more overlay planes for the display and composition of multiple layers of video or user interface elements. The display device 318 can be an internal or external display device. In one embodiment the display device 318 is a head mounted display device, such as a

virtual reality (VR) display device or an augmented reality (AR) display device. In some embodiments, graphics processor 300 includes a video codec engine 306 to encode, decode, or transcode media to, from, or between one or more media encoding formats, including, but not limited to Moving Picture Experts Group (MPEG) formats such as MPEG-2, Advanced Video Coding (AVC) formats such as H.264/MPEG-4 AVC, H.265/HEVC, Alliance for Open Media (AOMedia) VP8, VP9, AV1, AV2, as well as the Society of Motion Picture & Television Engineers (SMPTE) 421M/VC-1, and Joint Photographic Experts Group (JPEG) formats such as JPEG, and Motion JPEG (MJPEG) formats.

[0099] In some embodiments, graphics processor 300 includes a block image transfer (BLIT) engine to perform two-dimensional (2D) rasterizer operations including, for example, bit-boundary block transfers. However, in one embodiment, 2D graphics operations are performed using one or more components of graphics processing engine (GPE) 310. In some embodiments, GPE 310 is a compute engine for performing graphics operations, including three-dimensional (3D) graphics operations and media operations.

[0100] In some embodiments, GPE 310 includes a 3D pipeline 312 for performing 3D operations, such as rendering three-dimensional images and scenes using processing functions that act upon 3D primitive shapes (e.g., rectangle, triangle, etc.). The 3D pipeline 312 includes programmable and fixed function elements that perform various tasks within the element and/or spawn execution threads to a 3D/Media subsystem 315. While 3D pipeline 312 can be used to perform media operations, an embodiment of GPE 310 also includes a media pipeline 316 that is specifically used to perform media operations, such as video post-processing and image enhancement.

[0101] In some embodiments, media pipeline 316 includes fixed function or programmable logic units to perform one or more specialized media operations, such as video decode acceleration, video de-interlacing, and video encode acceleration in place of, or on behalf of video codec engine 306. In some embodiments, media pipeline 316 additionally includes a thread spawning unit to spawn threads for execution on 3D/Media subsystem 315. The spawned threads perform computations for the media operations on one or more graphics cores included in 3D/Media subsystem 315.

[0102] In some embodiments, 3D/Media subsystem 315 includes logic for executing threads spawned by 3D pipeline 312 and media pipeline 316. In one embodiment, the pipelines send thread execution requests to 3D/Media subsystem 315, which includes thread dispatch logic for arbitrating and dispatching the various requests to available thread execution resources. The execution resources include an array of graphics cores to process the 3D and media threads. In some embodiments, 3D/Media subsystem 315 includes one or more internal caches for thread instructions and data. In some embodiments, the subsystem also includes shared memory, including registers and addressable memory, to share data between threads and to store output data.

[0103] FIG. 3B illustrates a graphics processor 320 having a tiled architecture, according to embodiments described herein. In one embodiment the graphics processor 320 includes a graphics processing engine cluster 322 having multiple instances of the graphics processing engine 310 of FIG. 3A within a graphics engine tile 310A-310D. Each graphics engine tile 310A-310D can be interconnected via a set of tile interconnects 323A-323F. Each graphics engine

tile **310A-310D** can also be connected to a memory module or memory device **326A-326D** via memory interconnects **325A-325D**. The memory devices **326A-326D** can use any graphics memory technology. For example, the memory devices **326A-326D** may be graphics double data rate (GDDR) memory. The memory devices **326A-326D**, in one embodiment, are HBM modules that can be on-die with their respective graphics engine tile **310A-310D**. In one embodiment the memory devices **326A-326D** are stacked memory devices that can be stacked on top of their respective graphics engine tile **310A-310D**. In one embodiment, each graphics engine tile **310A-310D** and associated memory **326A-326D** reside on separate chiplets, which are bonded to a base die or base substrate, as described on further detail in FIGS. 11B-11D.

[0104] The graphics processor **320** may be configured with a non-uniform memory access (NUMA) system in which memory devices **326A-326D** are coupled with associated graphics engine tiles **310A-310D**. A given memory device may be accessed by graphics engine tiles other than the tile to which it is directly connected. However, access latency to the memory devices **326A-326D** may be lowest when accessing a local tile. In one embodiment, a cache coherent NUMA (ccNUMA) system is enabled that uses the tile interconnects **323A-323F** to enable communication between cache controllers within the graphics engine tiles **310A-310D** to maintain a consistent memory image when more than one cache stores the same memory location.

[0105] The graphics processing engine cluster **322** can connect with an on-chip or on-package fabric interconnect **324**. In one embodiment the fabric interconnect **324** includes a network processor, network on a chip (NoC), or another switching processor to enable the fabric interconnect **324** to act as a packet switched fabric interconnect that switches data packets between components of the graphics processor **320**. The fabric interconnect **324** can enable communication between graphics engine tiles **310A-310D** and components such as the video codec engine **306** and one or more copy engines **304**. The copy engines **304** can be used to move data out of, into, and between the memory devices **326A-326D** and memory that is external to the graphics processor **320** (e.g., system memory). The fabric interconnect **324** can also couple with one or more of the tile interconnects **323A-323F** to facilitate or enhance the interconnection between the graphics engine tiles **310A-310D**. The fabric interconnect **324** is also configurable to interconnect multiple instances of the graphics processor **320** (e.g., via the host interface **328**), enabling tile-to-tile communication between graphics engine tiles **310A-310D** of multiple GPUs. In one embodiment, the graphics engine tiles **310A-310D** of multiple GPUs can be presented to a host system as a single logical device.

[0106] The graphics processor **320** may optionally include a display controller **302** to enable a connection with the display device **318**. The graphics processor may also be configured as a graphics or compute accelerator. In the accelerator configuration, the display controller **302** and display device **318** may be omitted.

[0107] The graphics processor **320** can connect to a host system via a host interface **328**. The host interface **328** can enable communication between the graphics processor **320**, system memory, and/or other system components. The host interface **328** can be, for example a PCI express bus or another type of host system interface. For example, the host

interface **328** may be an NVLink or NVSwitch interface. The host interface **328** and fabric interconnect **324** can cooperate to enable multiple instances of the graphics processor **320** to act as single logical device. Cooperation between the host interface **328** and fabric interconnect **324** can also enable the individual graphics engine tiles **310A-310D** to be presented to the host system as distinct logical graphics devices.

[0108] FIG. 3C illustrates a compute accelerator **330**, according to embodiments described herein. The compute accelerator **330** can include architectural similarities with the graphics processor **320** of FIG. 3B and is optimized for compute acceleration. A compute engine cluster **332** can include a set of compute engine tiles **340A-340D** that include execution logic that is optimized for parallel or vector-based general-purpose compute operations. In some embodiments, the compute engine tiles **340A-340D** do not include fixed function graphics processing logic, although in one embodiment one or more of the compute engine tiles **340A-340D** can include logic to perform media acceleration. The compute engine tiles **340A-340D** can connect to memory **326A-326D** via memory interconnects **325A-325D**. The memory **326A-326D** and memory interconnects **325A-325D** may be similar technology as in graphics processor **320** or can be different. The compute engine tiles **340A-340D** can also be interconnected via a set of tile interconnects **323A-323F** and may be connected with and/or interconnected by a fabric interconnect **324**. Cross-tile communications can be facilitated via the fabric interconnect **324**. The fabric interconnect **324** (e.g., via the host interface **328**) can also facilitate communication between compute engine tiles **340A-340D** of multiple instances of the compute accelerator **330**. In one embodiment the compute accelerator **330** includes a large L3 cache **336** that can be configured as a device-wide cache. The compute accelerator **330** can also connect to a host processor and memory via a host interface **328** in a similar manner as the graphics processor **320** of FIG. 3B.

[0109] The compute accelerator **330** can also include an integrated network interface **342**. In one embodiment the network interface **342** includes a network processor and controller logic that enables the compute engine cluster **332** to communicate over a physical layer interconnect **344** without requiring data to traverse memory of a host system. In one embodiment, one of the compute engine tiles **340A-340D** is replaced by network processor logic and data to be transmitted or received via the physical layer interconnect **344** may be transmitted directly to or from memory **326A-326D**. Multiple instances of the compute accelerator **330** may be joined via the physical layer interconnect **344** into a single logical device. Alternatively, the various compute engine tiles **340A-340D** may be presented as distinct network accessible compute accelerator devices.

Graphics Processing Engine

[0110] FIG. 4 is a block diagram of a graphics processing engine **410** of a graphics processor in accordance with some embodiments. In one embodiment, the graphics processing engine (GPE) **410** is a version of the GPE **310** shown in FIG. 3A and may also represent a graphics engine tile **310A-310D** of FIG. 3B. Elements of FIG. 4 having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such. For

example, the 3D pipeline **312** and media pipeline **316** of FIG. 3A are illustrated. The media pipeline **316** is optional in some embodiments of the GPE **410** and may not be explicitly included within the GPE **410**. For example and in at least one embodiment, a separate media and/or image processor is coupled to the GPE **410**.

[0111] In some embodiments, GPE **410** couples with or includes a command streamer **403**, which provides a command stream to the 3D pipeline **312** and/or media pipelines **316**. Alternatively or additionally, the command streamer **403** may be directly coupled to a unified return buffer **418**. The unified return buffer **418** may be communicatively coupled to a graphics core cluster **414**. In some embodiments, command streamer **403** is coupled with memory, which can be system memory, or one or more of internal cache memory and shared cache memory. In some embodiments, command streamer **403** receives commands from the memory and sends the commands to 3D pipeline **312** and/or media pipeline **316**. The commands are directives fetched from a ring buffer, which stores commands for the 3D pipeline **312** and media pipeline **316**. In one embodiment, the ring buffer can additionally include batch command buffers storing batches of multiple commands. The commands for the 3D pipeline **312** can also include references to data stored in memory, such as but not limited to vertex and geometry data for the 3D pipeline **312** and/or image data and memory objects for the media pipeline **316**. The 3D pipeline **312** and media pipeline **316** process the commands and data by performing operations via logic within the respective pipelines or by dispatching one or more execution threads to a graphics core cluster **414**. In one embodiment the graphics core cluster **414** include one or more blocks of graphics cores (e.g., graphics core block **415A**, graphics core block **415B**), each block including one or more graphics cores. Each graphics core includes a set of graphics execution resources that includes general-purpose and graphics specific execution logic to perform graphics and compute operations, as well as fixed function texture processing and/or machine learning and artificial intelligence acceleration logic, such as matrix or AI acceleration logic.

[0112] In various embodiments the 3D pipeline **312** can include fixed function and programmable logic to process one or more shader programs, such as vertex shaders, geometry shaders, pixel shaders, fragment shaders, compute shaders, or other shader and/or GPGPU programs, by processing the instructions and dispatching execution threads to the graphics core cluster **414**. The graphics core cluster **414** provides a unified block of execution resources for use in processing these shader programs. Multi-purpose execution logic within the graphics core blocks **415A-415B** of the graphics core cluster **414** includes support for various 3D API shader languages and can execute multiple simultaneous execution threads associated with multiple shaders.

[0113] In some embodiments, the graphics core cluster **414** includes execution logic to perform media functions, such as video and/or image processing. In one embodiment, the graphics cores include general-purpose logic that is programmable to perform parallel general-purpose computational operations, in addition to graphics processing operations. The general-purpose logic can perform processing operations in parallel or in conjunction with general-purpose logic within the processor core(s) **107** of FIG. 1 or core **202A-202N** as in FIG. 2A.

[0114] Output data generated by threads executing on the graphics core cluster **414** can output data to memory in a unified return buffer (URB) **418**. The URB **418** can store data for multiple threads. In some embodiments the URB **418** may be used to send data between different threads executing on the graphics core cluster **414**. In some embodiments the URB **418** may additionally be used for synchronization between threads on the graphics core array and fixed function logic within the shared function logic **420**.

[0115] In some embodiments, graphics core cluster **414** is scalable, such that the cluster includes a variable number of graphics cores, each having a variable number of graphics cores based on the target power and performance level of GPE **410**. In one embodiment the execution resources are dynamically scalable, such that execution resources may be enabled or disabled as needed.

[0116] The graphics core cluster **414** couples with shared function logic **420** that includes multiple resources that are shared between the graphics cores in the graphics core array. The shared functions within the shared function logic **420** are hardware logic units that provide specialized supplemental functionality to the graphics core cluster **414**. In various embodiments, shared function logic **420** may include, but is not limited to sampler **421**, math **422**, and inter-thread communication (ITC) **423** logic. Additionally, some embodiments implement one or more cache(s) **425** within the shared function logic **420**. The shared function logic **420** can implement the same or similar functionality as the additional fixed function logic **238** of FIG. 2B.

[0117] A shared function is implemented at least in a case where the demand for a given specialized function is insufficient for inclusion within the graphics core cluster **414**. Instead, a single instantiation of that specialized function is implemented as a stand-alone entity in the shared function logic **420** and shared among the execution resources within the graphics core cluster **414**. The precise set of functions that are shared between the graphics core cluster **414** and included within the graphics core cluster **414** varies across embodiments. In some embodiments, specific shared functions within the shared function logic **420** that are used extensively by the graphics core cluster **414** may be included within shared function logic **416** within the graphics core cluster **414**. In various embodiments, the shared function logic **416** within the graphics core cluster **414** can include some or all logic within the shared function logic **420**. In one embodiment, all logic elements within the shared function logic **420** may be duplicated within the shared function logic **416** of the graphics core cluster **414**. In one embodiment the shared function logic **420** is excluded in favor of the shared function logic **416** within the graphics core cluster **414**.

Graphics Processing Resources

[0118] FIG. 5A-5C illustrate execution logic including an array of processing elements employed in a graphics processor, according to embodiments described herein. FIG. 5A illustrates graphics core cluster, according to an embodiment. FIG. 5B illustrates a vector engine of a graphics core, according to an embodiment. FIG. 5C illustrates a matrix engine of a graphics core, according to an embodiment. Elements of FIG. 5A-5C having the same reference numbers as the elements of any other figure herein may operate or function in any manner similar to that described elsewhere herein, but are not limited as such. For example, the elements of FIG. 5A-5C can be considered in the context of the

graphics processor core block 219 of FIG. 2B, and/or the graphics core blocks 415A-415B of FIG. 4. In one embodiment, the elements of FIG. 5A-5C have similar functionality to equivalent components of the graphics processor 208 of FIG. 2A, the GPU 239 of FIG. 2C or the GPGPU 270 of FIG. 2D.

[0119] As shown in FIG. 5A, in one embodiment the graphics core cluster 414 includes a graphics core block 415, which may be graphics core block 415A or graphics core block 415B of FIG. 4. The graphics core block 415 can include any number of graphics cores (e.g., graphics core 515A, graphics core 515B, through graphics core 515N). Multiple instances of the graphics core block 415 may be included. In one embodiment the elements of the graphics cores 515A-515N have similar or equivalent functionality as the elements of the graphics cores 221A-221F of FIG. 2B. In such embodiment, the graphics cores 515A-515N each include circuitry including but not limited to vector engines 502A-502N, matrix engines 503A-503N, memory load/store units 504A-504N, instruction caches 505A-505N, data caches/shared local memory 506A-506N, ray tracing units 508A-508N, samplers 510A-510N. The circuitry of the graphics cores 515A-515N can additionally include fixed function logic 512A-512N. The number of vector engines 502A-502N and matrix engines 503A-503N within the graphics cores 515A-515N of a design can vary based on the workload, performance, and power targets for the design.

[0120] With reference to graphics core 515A, the vector engine 502A and matrix engine 503A are configurable to perform parallel compute operations on data in a variety of integer and floating-point data formats based on instructions associated with shader programs. Each vector engine 502A and matrix engine 503A can act as a programmable general-purpose computational unit that is capable of executing multiple simultaneous hardware threads while processing multiple data elements in parallel for each thread. The vector engine 502A and matrix engine 503A support the processing of variable width vectors at various SIMD widths, including but not limited to SIMD8, SIMD16, and SIMD32. Input data elements can be stored as a packed data type in a register and the vector engine 502A and matrix engine 503A can process the various elements based on the data size of the elements. For example, when operating on a 256-bit wide vector, the 256 bits of the vector are stored in a register and the vector is processed as four separate 64-bit packed data elements (Quad-Word (QW) size data elements), eight separate 32-bit packed data elements (Double Word (DW) size data elements), sixteen separate 16-bit packed data elements (Word (W) size data elements), or thirty-two separate 8-bit data elements (byte (B) size data elements). However, different vector widths and register sizes are possible. In one embodiment, the vector engine 502A and matrix engine 503A are also configurable for SIMT operation on warps or thread groups of various sizes (e.g., 8, 16, or 32 threads).

[0121] Continuing with graphics core 515A, the memory load/store unit 504A services memory access requests that are issued by the vector engine 502A, matrix engine 503A, and/or other components of the graphics core 515A that have access to memory. The memory access request can be processed by the memory load/store unit 504A to load or store the requested data to or from cache or memory into a register file associated with the vector engine 502A and/or matrix engine 503A. The memory load/store unit 504A can also perform prefetching operations. In one embodiment, the

memory load/store unit 504A is configured to provide SIMT scatter/gather prefetching or block prefetching for data stored in memory 610 of FIG. 6, from memory that is local to other tiles via the tile interconnect 608 of FIG. 6, or from system memory. Referring again to FIG. 6, prefetching can be performed to a specific L1 cache (e.g., data cache/shared local memory 506A), the L2 cache 604 or the L3 cache 606. In one embodiment, a prefetch to the L3 cache 606 automatically results in the data being stored in the L2 cache 604.

[0122] Referring to FIG. 5A, the instruction cache 505A stores instructions to be executed by the graphics core 515A. In one embodiment, the graphics core 515A also includes instruction fetch and prefetch circuitry that fetches or prefetches instructions into the instruction cache 505A. The graphics core 515A also includes instruction decode logic to decode instructions within the instruction cache 505A. The data cache/shared local memory 506A can be configured as a data cache that is managed by a cache controller that implements a cache replacement policy and/or configured as explicitly managed shared memory. The ray tracing unit 508A includes circuitry to accelerate ray tracing operations. The sampler 510A provides texture sampling for 3D operations and media sampling for media operations. The fixed function logic 512A includes fixed function circuitry that is shared between the various instances of the vector engine 502A and matrix engine 503A. Graphics cores 515B-515N can operate in a similar manner as graphics core 515A.

[0123] Functionality of the instruction caches 505A-505N, data caches/shared local memory 506A-506N, ray tracing units 508A-508N, samplers 510A-2710N, and fixed function logic 512A-512N corresponds with equivalent functionality in the graphics processor architectures described herein. For example, the instruction caches 505A-505N can operate in a similar manner as instruction cache 255 of FIG. 2D. The data caches/shared local memory 506A-506N, ray tracing units 508A-508N, and samplers 510A-2710N can operate in a similar manner as the cache/SLM 228A-228F, ray tracing units 227A-227F, and samplers 226A-226F of FIG. 2B. The fixed function logic 512A-512N can include elements of the geometry/fixed function pipeline 231 and/or additional fixed function logic 238 of FIG. 2B. In one embodiment, the ray tracing units 508A-508N include circuitry to perform ray tracing acceleration operations performed by the ray tracing cores 245 of FIG. 2C.

[0124] As shown in FIG. 5B, in one embodiment the vector engine 502 includes an instruction fetch unit 537, a general register file array (GRF) 524, an architectural register file array (ARF) 526, a thread arbiter 522, a send unit 530, a branch unit 532, a set of SIMD floating point units (FPUs) 534, and in one embodiment a set of integer SIMD ALUs 535. The GRF 524 and ARF 526 includes the set of general register files and architecture register files associated with each hardware thread that may be active in the vector engine 502. In one embodiment, per thread architectural state is maintained in the ARF 526, while data used during thread execution is stored in the GRF 524. The execution state of each thread, including the instruction pointers for each thread, can be held in thread-specific registers in the ARF 526.

[0125] In one embodiment the vector engine 502 has an architecture that is a combination of Simultaneous Multi-Threading (SMT) and fine-grained Interleaved Multi-Threading (IMT). The architecture has a modular configu-

ration that can be fine-tuned at design time based on a target number of simultaneous threads and number of registers per graphics core, where graphics core resources are divided across logic used to execute multiple simultaneous threads. The number of logical threads that may be executed by the vector engine **502** is not limited to the number of hardware threads, and multiple logical threads can be assigned to each hardware thread.

[0126] In one embodiment, the vector engine **502** can co-issue multiple instructions, which may each be different instructions. The thread arbiter **522** can dispatch the instructions to one of the send unit **530**, branch unit **532**, or SIMD FPU(s) **534** for execution. Each execution thread can access 128 general-purpose registers within the GRF **524**, where each register can store 32 bytes, accessible as a variable width vector of 32-bit data elements. In one embodiment, each thread has access to 4 Kbytes within the GRF **524**, although embodiments are not so limited, and greater or fewer register resources may be provided in other embodiments. In one embodiment the vector engine **502** is partitioned into seven hardware threads that can independently perform computational operations, although the number of threads per vector engine **502** can also vary according to embodiments. For example, in one embodiment up to 16 hardware threads are supported. In an embodiment in which seven threads may access 4 Kbytes, the GRF **524** can store a total of 28 Kbytes. Where 16 threads may access 4 Kbytes, the GRF **524** can store a total of 64 Kbytes. Flexible addressing modes can permit registers to be addressed together to build effectively wider registers or to represent strided rectangular block data structures.

[0127] In one embodiment, memory operations, sampler operations, and other longer-latency system communications are dispatched via “send” instructions that are executed by the message passing send unit **530**. In one embodiment, branch instructions are dispatched to a dedicated branch unit **532** to facilitate SIMD divergence and eventual convergence.

[0128] In one embodiment the vector engine **502** includes one or more SIMD floating point units (FPU(s)) **534** to perform floating-point operations. In one embodiment, the FPU(s) **534** also support integer computation. In one embodiment the FPU(s) **534** can execute up to M number of 32-bit floating-point (or integer) operations, or execute up to 2M 16-bit integer or 16-bit floating-point operations. In one embodiment, at least one of the FPU(s) provides extended math capability to support high-throughput transcendental math functions and double precision 64-bit floating-point. In some embodiments, a set of 8-bit integer SIMD ALUs **535** are also present and may be specifically optimized to perform operations associated with machine learning computations. In one embodiment, the SIMD ALUs are replaced by an additional set of SIMD FPUs **534** that are configurable to perform integer and floating-point operations. In one embodiment, the SIMD FPUs **534** and SIMD ALUs **535** are configurable to execute SIMT programs. In one embodiment, combined SIMD+SIMT operation is supported.

[0129] In one embodiment, arrays of multiple instances of the vector engine **502** can be instantiated in a graphics core. For scalability, product architects can choose the exact number of vector engines per graphics core grouping. In one embodiment the vector engine **502** can execute instructions across a plurality of execution channels. In a further embodiment,

each thread executed on the vector engine **502** is executed on a different channel.

[0130] As shown in FIG. 5C, in one embodiment the matrix engine **503** includes an array of processing elements that are configured to perform tensor operations including vector/matrix and matrix/matrix operations, such as but not limited to matrix multiply and/or dot product operations. The matrix engine **503** is configured with M rows and N columns of processing elements (**552AA-552MN**) that include multiplier and adder circuits organized in a pipelined fashion. In one embodiment, the processing elements **552AA-552MN** make up the physical pipeline stages of an N wide and M deep systolic array that can be used to perform vector/matrix or matrix/matrix operations in a data-parallel manner, including matrix multiply, fused multiply-add, dot product or other general matrix-matrix multiplication (GEMM) operations. In one embodiment the matrix engine **503** supports 16-bit floating point operations, as well as 8-bit, 4-bit, 2-bit, and binary integer operations.

[0131] The matrix engine **503** can also be configured to accelerate specific machine learning operations. In such embodiments, the matrix engine **503** can be configured with support for the bfloat (brain floating point) 16-bit floating point format or a tensor float 32-bit floating point format (TF32) that have different numbers of mantissa and exponent bits relative to Institute of Electrical and Electronics Engineers (IEEE) 754 formats.

[0132] In one embodiment, during each cycle, each stage can add the result of operations performed at that stage to the output of the previous stage. In other embodiments, the pattern of data movement between the processing elements **552AA-552MN** after a set of computational cycles can vary based on the instruction or macro-operation being performed. For example, in one embodiment partial sum loopback is enabled and the processing elements may instead add the output of a current cycle with output generated in the previous cycle. In one embodiment, the final stage of the systolic array can be configured with a loopback to the initial stage of the systolic array. In such embodiment, the number of physical pipeline stages may be decoupled from the number of logical pipeline stages that are supported by the matrix engine **503**. For example, where the processing elements **552AA-552MN** are configured as a systolic array of M physical stages, a loopback from stage M to the initial pipeline stage can enable the processing elements **552AA-552MN** to operate as a systolic array of, for example, 2M, 3M, 4M, etc., logical pipeline stages.

[0133] In one embodiment, the matrix engine **503** includes memory **541A-541N**, **542A-542M** to store input data in the form of row and column data for input matrices. Memory **542A-542M** is configurable to store row elements (A0-Am) of a first input matrix and memory **541A-541N** is configurable to store column elements (B0-Bn) of a second input matrix. The row and column elements are provided as input to the processing elements **552AA-552MN** for processing. In one embodiment, row and column elements of the input matrices can be stored in a systolic register file **540** within the matrix engine **503** before those elements are provided to the memory **541A-541N**, **542A-542M**. In one embodiment, the systolic register file **540** is excluded and the memory **541A-541N**, **542A-542M** is loaded from registers in an associated vector engine (e.g., GRF **524** of vector engine **502** of FIG. 5B) or other memory of the graphics core that includes the matrix engine **503** (e.g., data cache/shared local

memory **506A** for matrix engine **503A** of FIG. 5A). Results generated by the processing elements **552AA-552MN** are then output to an output buffer and/or written to a register file (e.g., systolic register file **540**, GRF **524**, data cache/shared local memory **506A-506N**) for further processing by other functional units of the graphics processor or for output to memory.

[0134] In some embodiments, the matrix engine **503** is configured with support for input sparsity, where multiplication operations for sparse regions of input data can be bypassed by skipping multiply operations that have a zero-value operand. In one embodiment, the processing elements **552AA-552MN** are configured to skip the performance of certain operations that have zero value input. In one embodiment, sparsity within input matrices can be detected and operations having known zero output values can be bypassed before being submitted to the processing elements **552AA-552MN**. The loading of zero value operands into the processing elements can be bypassed and the processing elements **552AA-552MN** can be configured to perform multiplications on the non-zero value input elements. The matrix engine **503** can also be configured with support for output sparsity, such that operations with results that are pre-determined to be zero are bypassed. For input sparsity and/or output sparsity, in one embodiment, metadata is provided to the processing elements **552AA-552MN** to indicate, for a processing cycle, which processing elements and/or data channels are to be active during that cycle.

[0135] In one embodiment, the matrix engine **503** includes hardware to enable operations on sparse data having a compressed representation of a sparse matrix that stores non-zero values and metadata that identifies the positions of the non-zero values within the matrix. Exemplary compressed representations include but are not limited to compressed tensor representations such as compressed sparse row (CSR), compressed sparse column (CSC), compressed sparse fiber (CSF) representations. Support for compressed representations enable operations to be performed on input in a compressed tensor format without requiring the compressed representation to be decompressed or decoded. In such embodiment, operations can be performed only on non-zero input values and the resulting non-zero output values can be mapped into an output matrix. In some embodiments, hardware support is also provided for machine-specific lossless data compression formats that are used when transmitting data within hardware or across system busses. Such data may be retained in a compressed format for sparse input data and the matrix engine **503** can use the compression metadata for the compressed data to enable operations to be performed on only non-zero values, or to enable blocks of zero data input to be bypassed for multiply operations.

[0136] In various embodiments, input data can be provided by a programmer in a compressed tensor representation, or a codec can compress input data into the compressed tensor representation or another sparse data encoding. In addition to support for compressed tensor representations, streaming compression of sparse input data can be performed before the data is provided to the processing elements **552AA-552MN**. In one embodiment, compression is performed on data written to a cache memory associated with the graphics core cluster **414**, with the compression being performed with an encoding that is supported by the matrix engine **503**. In one embodiment, the matrix engine

503 includes support for input having structured sparsity in which a pre-determined level or pattern of sparsity is imposed on input data. This data may be compressed to a known compression ratio, with the compressed data being processed by the processing elements **552AA-552MN** according to metadata associated with the compressed data.

[0137] FIG. 6 illustrates a tile **600** of a multi-tile processor, according to an embodiment. In one embodiment, the tile **600** is representative of one of the graphics engine tiles **310A-310D** of FIG. 3B or compute engine tiles **340A-340D** of FIG. 3C. The tile **600** of the multi-tile graphics processor includes an array of graphics core clusters (e.g., graphics core cluster **414A**, graphics core cluster **414B**, through graphics core cluster **414N**), with each graphics core cluster having an array of graphics cores **515A-515N**. The tile **600** also includes a global dispatcher **602** to dispatch threads to processing resources of the tile **600**.

[0138] The tile **600** can include or couple with an L3 cache **606** and memory **610**. In various embodiments, the L3 cache **606** may be excluded or the tile **600** can include additional levels of cache, such as an LA cache. In one embodiment, each instance of the tile **600** in the multi-tile graphics processor has an associated memory **610**, such as in FIG. 3B and FIG. 3C. In one embodiment, a multi-tile processor can be configured as a multi-chip module in which the L3 cache **606** and/or memory **610** reside on separate chiplets than the graphics core clusters **414A-414N**. In this context, a chiplet is an at least partially packaged integrated circuit that includes distinct units of logic that can be assembled with other chiplets into a larger package. For example, the L3 cache **606** can be included in a dedicated cache chiplet or can reside on the same chiplet as the graphics core clusters **414A-414N**. In one embodiment, the L3 cache **606** can be included in an active base die or active interposer, as illustrated in FIG. 11C.

[0139] A memory fabric **603** enables communication among the graphics core clusters **414A-414N**, L3 cache **606**, and memory **610**. An L2 cache **604** couples with the memory fabric **603** and is configurable to cache transactions performed via the memory fabric **603**. A tile interconnect **608** enables communication with other tiles on the graphics processors and may be one of tile interconnects **323A-323F** of FIGS. 3B and 3C. In embodiments in which the L3 cache **606** is excluded from the tile **600**, the L2 cache **604** may be configured as a combined L2/L3 cache. The memory fabric **603** is configurable to route data to the L3 cache **606** or memory controllers associated with the memory **610** based on the presence or absence of the L3 cache **606** in a specific implementation. The L3 cache **606** can be configured as a per-tile cache that is dedicated to processing resources of the tile **600** or may be a partition of a GPU-wide L3 cache.

[0140] FIG. 7 is a block diagram illustrating graphics processor instruction formats **700** according to some embodiments. In one or more embodiment, the graphics processor cores support an instruction set having instructions in multiple formats. The solid lined boxes illustrate the components that are generally included in a graphics core instruction, while the dashed lines include components that are optional or that are only included in a sub-set of the instructions. In some embodiments, the graphics processor instruction format **700** described and illustrated are macro-instructions, in that they are instructions supplied to the graphics core, as opposed to micro-operations resulting from

instruction decode once the instruction is processed. Thus, a single instruction may cause hardware to perform multiple micro-operations.

[0141] In some embodiments, the graphics processor natively supports instructions in a 128-bit instruction format **710**. A 64-bit compacted instruction format **730** is available for some instructions based on the selected instruction, instruction options, and number of operands. The native 128-bit instruction format **710** provides access to all instruction options, while some options and operations are restricted in the 64-bit format **730**. The native instructions available in the 64-bit format **730** vary by embodiment. In some embodiments, the instruction is compacted in part using a set of index values in an index field **713**. The graphics core hardware references a set of compaction tables based on the index values and uses the compaction table outputs to reconstruct a native instruction in the 128-bit instruction format **710**. Other sizes and formats of instruction can be used.

[0142] For each format, instruction opcode **712** defines the operation that the graphics core is to perform. The graphics cores execute each instruction in parallel across the multiple data elements of each operand. For example, in response to an add instruction the graphics core performs a simultaneous add operation across each color channel representing a texture element or picture element. By default, the graphics core performs each instruction across all data channels of the operands. In some embodiments, instruction control field **714** enables control over certain execution options, such as channels selection (e.g., predication) and data channel order (e.g., swizzle). For instructions in the 128-bit instruction format **710** an exec-size field **716** limits the number of data channels that will be executed in parallel. In some embodiments, exec-size field **716** is not available for use in the 64-bit compact instruction format **730**.

[0143] Some graphics core instructions have up to three operands including two source operands, **src0 720**, **src1 722**, and one destination **718**. In some embodiments, the graphics cores support dual destination instructions, where one of the destinations is implied. Data manipulation instructions can have a third source operand (e.g., **SRC2 724**), where the instruction opcode **712** determines the number of source operands. An instruction's last source operand can be an immediate (e.g., hard-coded) value passed with the instruction.

[0144] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726** specifying, for example, whether direct register addressing mode or indirect register addressing mode is used. When direct register addressing mode is used, the register address of one or more operands is directly provided by bits in the instruction.

[0145] In some embodiments, the 128-bit instruction format **710** includes an access/address mode field **726**, which specifies an address mode and/or an access mode for the instruction. In one embodiment the access mode is used to define a data access alignment for the instruction. Some embodiments support access modes including a 16-byte aligned access mode and a 1-byte aligned access mode, where the byte alignment of the access mode determines the access alignment of the instruction operands. For example, when in a first mode, the instruction may use byte-aligned addressing for source and destination operands and when in

a second mode, the instruction may use 16-byte-aligned addressing for all source and destination operands.

[0146] In one embodiment, the address mode portion of the access/address mode field **726** determines whether the instruction is to use direct or indirect addressing. When direct register addressing mode is used bits in the instruction directly provide the register address of one or more operands. When indirect register addressing mode is used, the register address of one or more operands may be computed based on an address register value and an address immediate field in the instruction.

[0147] In some embodiments instructions are grouped based on opcode **712** bit-fields to simplify Opcode decode **740**. For an 8-bit opcode, bits 4, 5, and 6 allow the graphics core to determine the type of opcode. The precise opcode grouping shown is merely an example. In some embodiments, a move and logic opcode group **742** includes data movement and logic instructions (e.g., move (mov), compare (cmp)). In some embodiments, move and logic group **742** shares the five most significant bits (MSB), where move (mov) instructions are in the form of 0000xxxxb and logic instructions are in the form of 0001xxxxb. A flow control instruction group **744** (e.g., call, jump (jmp)) includes instructions in the form of 0010xxxxb (e.g., 0x20). A miscellaneous instruction group **746** includes a mix of instructions, including synchronization instructions (e.g., wait, send) in the form of 0011xxxxb (e.g., 0x30). A parallel math instruction group **748** includes component-wise arithmetic instructions (e.g., add, multiply (mul)) in the form of 0100xxxxb (e.g., 0x40). The parallel math instruction group **748** performs the arithmetic operations in parallel across data channels. The vector math group **750** includes arithmetic instructions (e.g., dp4) in the form of 0101xxxxb (e.g., 0x50). The vector math group performs arithmetic such as dot product calculations on vector operands. The illustrated opcode decode **740**, in one embodiment, can be used to determine which portion of a graphics core will be used to execute a decoded instruction. For example, some instructions may be designated as systolic instructions that will be performed by a systolic array. Other instructions, such as ray-tracing instructions (not shown) can be routed to a ray-tracing core or ray-tracing logic within a slice or partition of execution logic.

Graphics Pipeline

[0148] FIG. 8 is a block diagram of another embodiment of a graphics processor **800**. Elements of FIG. 8 having the same reference numbers (or names) as the elements of any other figure herein can operate or function in any manner similar to that described elsewhere herein, but are not limited to such.

[0149] In some embodiments, graphics processor **800** includes a geometry pipeline **820**, a media pipeline **830**, a display engine **840**, thread execution logic **850**, and a render output pipeline **870**. In some embodiments, graphics processor **800** is a graphics processor within a multi-core processing system that includes one or more general-purpose processing cores. The graphics processor is controlled by register writes to one or more control registers (not shown) or via commands issued to graphics processor **800** via a ring interconnect **802**. In some embodiments, ring interconnect **802** couples graphics processor **800** to other processing components, such as other graphics processors or general-purpose processors. Commands from ring intercon-

nect **802** are interpreted by a command streamer **803**, which supplies instructions to individual components of the geometry pipeline **820** or the media pipeline **830**.

[0150] In some embodiments, command streamer **803** directs the operation of a vertex fetcher **805** that reads vertex data from memory and executes vertex-processing commands provided by command streamer **803**. In some embodiments, vertex fetcher **805** provides vertex data to a vertex shader **807**, which performs coordinate space transformation and lighting operations to each vertex. In some embodiments, vertex fetcher **805** and vertex shader **807** execute vertex-processing instructions by dispatching execution threads to graphics cores **852A-852B** via a thread dispatcher **831**.

[0151] In some embodiments, graphics cores **852A-852B** are an array of vector processors having an instruction set for performing graphics and media operations. In some embodiments, graphics cores **852A-852B** have an attached L1 cache **851** that is specific for each array or shared between the arrays. The cache can be configured as a data cache, an instruction cache, or a single cache that is partitioned to contain data and instructions in different partitions.

[0152] In some embodiments, geometry pipeline **820** includes tessellation components to perform hardware-accelerated tessellation of 3D objects. In some embodiments, a programmable hull shader **811** configures the tessellation operations. A programmable domain shader **817** provides back-end evaluation of tessellation output. A tessellator **813** operates at the direction of hull shader **811** and contains special purpose logic to generate a set of detailed geometric objects based on a coarse geometric model that is provided as input to geometry pipeline **820**. In some embodiments, if tessellation is not used, tessellation components (e.g., hull shader **811**, tessellator **813**, and domain shader **817**) can be bypassed. The tessellation components can operate based on data received from the vertex shader **807**.

[0153] In some embodiments, complete geometric objects can be processed by a geometry shader **819** via one or more threads dispatched to graphics cores **852A-852B** or can proceed directly to the clipper **829**. In some embodiments, the geometry shader operates on entire geometric objects, rather than vertices or patches of vertices as in previous stages of the graphics pipeline. If the tessellation is disabled the geometry shader **819** receives input from the vertex shader **807**. In some embodiments, geometry shader **819** is programmable by a geometry shader program to perform geometry tessellation if the tessellation units are disabled.

[0154] Before rasterization, a clipper **829** processes vertex data. The clipper **829** may be a fixed function clipper or a programmable clipper having clipping and geometry shader functions. In some embodiments, a rasterizer and depth test component **873** in the render output pipeline **870** dispatches pixel shaders to convert the geometric objects into per pixel representations. In some embodiments, pixel shader logic is included in thread execution logic **850**. In some embodiments, an application can bypass the rasterizer and depth test component **873** and access un-rasterized vertex data via a stream out unit **823**.

[0155] The graphics processor **800** has an interconnect bus, interconnect fabric, or some other interconnect mechanism that allows data and message passing amongst the major components of the processor. In some embodiments, graphics cores **852A-852B** and associated logic units (e.g., L1 cache **851**, sampler **854**, texture cache **858**, etc.) inter-

connect via a data port **856** to perform memory access and communicate with render output pipeline components of the processor. In some embodiments, sampler **854**, caches **851**, **858** and graphics cores **852A-852B** each have separate memory access paths. In one embodiment the texture cache **858** can also be configured as a sampler cache.

[0156] In some embodiments, render output pipeline **870** contains a rasterizer and depth test component **873** that converts vertex-based objects into an associated pixel-based representation. In some embodiments, the rasterizer logic includes a windower/masker unit to perform fixed function triangle and line rasterization. An associated render cache **878** and depth cache **879** are also available in some embodiments. A pixel operations component **877** performs pixel-based operations on the data, though in some instances, pixel operations associated with 2D operations (e.g., bit block image transfers with blending) are performed by the 2D engine **841**, or substituted at display time by the display controller **843** using overlay display planes. In some embodiments, a shared L3 cache **875** is available to all graphics components, allowing the sharing of data without the use of main system memory.

[0157] In some embodiments, media pipeline **830** includes a media engine **837** and a video front-end **834**. In some embodiments, video front-end **834** receives pipeline commands from the command streamer **803**. In some embodiments, media pipeline **830** includes a separate command streamer. In some embodiments, video front-end **834** processes media commands before sending the command to the media engine **837**. In some embodiments, media engine **837** includes thread spawning functionality to spawn threads for dispatch to thread execution logic **850** via thread dispatcher **831**.

[0158] In some embodiments, graphics processor **800** includes a display engine **840**. In some embodiments, display engine **840** is external to processor **800** and couples with the graphics processor via the ring interconnect **802**, or some other interconnect bus or fabric. In some embodiments, display engine **840** includes a 2D engine **841** and a display controller **843**. In some embodiments, display engine **840** contains special purpose logic capable of operating independently of the 3D pipeline. In some embodiments, display controller **843** couples with a display device (not shown), which may be a system integrated display device, as in a laptop computer, or an external display device attached via a display device connector.

[0159] In some embodiments, the geometry pipeline **820** and media pipeline **830** are configurable to perform operations based on multiple graphics and media programming interfaces and are not specific to any one application programming interface (API). In some embodiments, driver software for the graphics processor translates API calls that are specific to a particular graphics or media library into commands that can be processed by the graphics processor. In some embodiments, support is provided for the Open Graphics Library (OpenGL), Open Computing Language (OpenCL), and/or Vulkan graphics and compute API, all from the Khronos Group. In some embodiments, support may also be provided for the Direct3D library from the Microsoft Corporation. In some embodiments, a combination of these libraries may be supported. Support may also be provided for the Open Source Computer Vision Library (OpenCV). A future API with a compatible 3D pipeline

would also be supported if a mapping can be made from the pipeline of the future API to the pipeline of the graphics processor.

Graphics Pipeline Programming

[0160] FIG. 9A is a block diagram illustrating a graphics processor command format 900 that may be used to program graphics processing pipelines according to some embodiments. FIG. 9B is a block diagram illustrating a graphics processor command sequence 910 according to an embodiment. The solid lined boxes in FIG. 9A illustrate the components that are generally included in a graphics command while the dashed lines include components that are optional or that are only included in a sub-set of the graphics commands. The exemplary graphics processor command format 900 of FIG. 9A includes data fields to identify a client 902, a command operation code (opcode) 904, and a data field 906 for the command. A sub-opcode 905 and a command size 908 are also included in some commands.

[0161] In some embodiments, client 902 specifies the client unit of the graphics device that processes the command data. In some embodiments, a graphics processor command parser examines the client field of each command to condition the further processing of the command and route the command data to the appropriate client unit. In some embodiments, the graphics processor client units include a memory interface unit, a render unit, a 2D unit, a 3D unit, and a media unit. Each client unit has a corresponding processing pipeline that processes the commands. Once the command is received by the client unit, the client unit reads the opcode 904 and, if present, sub-opcode 905 to determine the operation to perform. The client unit performs the command using information in data field 906. For some commands an explicit command size 908 is expected to specify the size of the command. In some embodiments, the command parser automatically determines the size of at least some of the commands based on the command opcode. In some embodiments commands are aligned via multiples of a double word. Other command formats can be used.

[0162] The flow diagram in FIG. 9B illustrates an exemplary graphics processor command sequence 910. In some embodiments, software or firmware of a data processing system that features an embodiment of a graphics processor uses a version of the command sequence shown to set up, execute, and terminate a set of graphics operations. A sample command sequence is shown and described for purposes of example only as embodiments are not limited to these specific commands or to this command sequence. Moreover, the commands may be issued as batch of commands in a command sequence, such that the graphics processor will process the sequence of commands in at least partially concurrence.

[0163] In some embodiments, the graphics processor command sequence 910 may begin with a pipeline flush command 912 to cause any active graphics pipeline to complete the currently pending commands for the pipeline. In some embodiments, the 3D pipeline 922 and the media pipeline 924 do not operate concurrently. The pipeline flush is performed to cause the active graphics pipeline to complete any pending commands. In response to a pipeline flush, the command parser for the graphics processor will pause command processing until the active drawing engines complete pending operations and the relevant read caches are invalidated. Optionally, any data in the render cache that is

marked ‘dirty’ can be flushed to memory. In some embodiments, pipeline flush command 912 can be used for pipeline synchronization or before placing the graphics processor into a low power state.

[0164] In some embodiments, a pipeline select command 913 is used when a command sequence requires the graphics processor to explicitly switch between pipelines. In some embodiments, a pipeline select command 913 is required only once within an execution context before issuing pipeline commands unless the context is to issue commands for both pipelines. In some embodiments, a pipeline flush command 912 is required immediately before a pipeline switch via the pipeline select command 913.

[0165] In some embodiments, a pipeline control command 914 configures a graphics pipeline for operation and is used to program the 3D pipeline 922 and the media pipeline 924. In some embodiments, pipeline control command 914 configures the pipeline state for the active pipeline. In one embodiment, the pipeline control command 914 is used for pipeline synchronization and to clear data from one or more cache memories within the active pipeline before processing a batch of commands.

[0166] In some embodiments, commands related to the return buffer state 916 are used to configure a set of return buffers for the respective pipelines to write data. Some pipeline operations require the allocation, selection, or configuration of one or more return buffers into which the operations write intermediate data during processing. In some embodiments, the graphics processor also uses one or more return buffers to store output data and to perform cross thread communication. In some embodiments, the return buffer state 916 includes selecting the size and number of return buffers to use for a set of pipeline operations.

[0167] The remaining commands in the command sequence differ based on the active pipeline for operations. Based on a pipeline determination 920, the command sequence is tailored to the 3D pipeline 922 beginning with the 3D pipeline state 930 or the media pipeline 924 beginning at the media pipeline state 940.

[0168] The commands to configure the 3D pipeline state 930 include 3D state setting commands for vertex buffer state, vertex element state, constant color state, depth buffer state, and other state variables that are to be configured before 3D primitive commands are processed. The values of these commands are determined at least in part based on the particular 3D API in use. In some embodiments, 3D pipeline state 930 commands are also able to selectively disable or bypass certain pipeline elements if those elements will not be used.

[0169] In some embodiments, 3D primitive 932 command is used to submit 3D primitives to be processed by the 3D pipeline. Commands and associated parameters that are passed to the graphics processor via the 3D primitive 932 command are forwarded to the vertex fetch function in the graphics pipeline. The vertex fetch function uses the 3D primitive 932 command data to generate vertex data structures. The vertex data structures are stored in one or more return buffers. In some embodiments, 3D primitive 932 command is used to perform vertex operations on 3D primitives via vertex shaders. To process vertex shaders, 3D pipeline 922 dispatches shader programs to the graphics cores.

[0170] In some embodiments, 3D pipeline 922 is triggered via an execute 934 command or event. In some embodi-

ments, a register write triggers command execution. In some embodiments execution is triggered via a ‘go’ or ‘kick’ command in the command sequence. In one embodiment, command execution is triggered using a pipeline synchronization command to flush the command sequence through the graphics pipeline. The 3D pipeline will perform geometry processing for the 3D primitives. Once operations are complete, the resulting geometric objects are rasterized and the pixel engine colors the resulting pixels. Additional commands to control pixel shading and pixel back-end operations may also be included for those operations.

[0171] In some embodiments, the graphics processor command sequence **910** follows the media pipeline **924** path when performing media operations. In general, the specific use and manner of programming for the media pipeline **924** depends on the media or compute operations to be performed. Specific media decode operations may be offloaded to the media pipeline during media decode. In some embodiments, the media pipeline can also be bypassed and media decode can be performed in whole or in part using resources provided by one or more general-purpose processing cores. In one embodiment, the media pipeline also includes elements for general-purpose graphics processor unit (GPGPU) operations, where the graphics processor is used to perform SIMD vector operations using computational shader programs that are not explicitly related to the rendering of graphics primitives.

[0172] In some embodiments, media pipeline **924** is configured in a similar manner as the 3D pipeline **922**. A set of commands to configure the media pipeline state **940** are dispatched or placed into a command queue before the media object commands **942**. In some embodiments, commands for the media pipeline state **940** include data to configure the media pipeline elements that will be used to process the media objects. This includes data to configure the video decode and video encode logic within the media pipeline, such as encode or decode format. In some embodiments, commands for the media pipeline state **940** also support the use of one or more pointers to “indirect” state elements that contain a batch of state settings.

[0173] In some embodiments, media object commands **942** supply pointers to media objects for processing by the media pipeline. The media objects include memory buffers containing video data to be processed. In some embodiments, all media pipeline states must be valid before issuing a media object command **942**. Once the pipeline state is configured and media object commands **942** are queued, the media pipeline **924** is triggered via an execute command **944** or an equivalent execute event (e.g., register write). Output from media pipeline **924** may then be post processed by operations provided by the 3D pipeline **922** or the media pipeline **924**. In some embodiments, GPGPU operations are configured and executed in a similar manner as media operations.

Graphics Software Architecture

[0174] FIG. 10 illustrates an exemplary graphics software architecture for a data processing system **1000** according to some embodiments. In some embodiments, software architecture includes a 3D graphics application **1010**, an operating system **1020**, and at least one processor **1030**. In some embodiments, processor **1030** includes a graphics processor **1032** and one or more general-purpose processor core(s)

1034. The graphics application **1010** and operating system **1020** each execute in the system memory **1050** of the data processing system.

[0175] In some embodiments, 3D graphics application **1010** contains one or more shader programs including shader instructions **1012**. The shader language instructions may be in a high-level shader language, such as the High-Level Shader Language (HLSL) of Direct3D, the OpenGL Shader Language (GLSL), and so forth. The application also includes executable instructions **1014** in a machine language suitable for execution by the general-purpose processor core **1034**. The application also includes graphics objects **1016** defined by vertex data.

[0176] In some embodiments, operating system **1020** is a Microsoft® Windows® operating system from the Microsoft Corporation, a proprietary UNIX-like operating system, or an open source UNIX-like operating system using a variant of the Linux kernel. The operating system **1020** can support a graphics API **1022** such as the Direct3D API, the OpenGL API, or the Vulkan API. When the Direct3D API is in use, the operating system **1020** uses a front-end shader compiler **1024** to compile any shader instructions **1012** in HLSL into a lower-level shader language. The compilation may be a just-in-time (JIT) compilation or the application can perform shader pre-compilation. In some embodiments, high-level shaders are compiled into low-level shaders during the compilation of the 3D graphics application **1010**. In some embodiments, the shader instructions **1012** are provided in an intermediate form, such as a version of the Standard Portable Intermediate Representation (SPIR) used by the Vulkan API.

[0177] In some embodiments, user mode graphics driver **1026** contains a back-end shader compiler **1027** to convert the shader instructions **1012** into a hardware specific representation. When the OpenGL API is in use, shader instructions **1012** in the GLSL high-level language are passed to a user mode graphics driver **1026** for compilation. In some embodiments, user mode graphics driver **1026** uses operating system kernel mode functions **1028** to communicate with a kernel mode graphics driver **1029**. In some embodiments, kernel mode graphics driver **1029** communicates with graphics processor **1032** to dispatch commands and instructions.

IP Core Implementations

[0178] One or more aspects of at least one embodiment may be implemented by representative code stored on a machine-readable medium which represents and/or defines logic within an integrated circuit such as a processor. For example, the machine-readable medium may include instructions which represent various logic within the processor. When read by a machine, the instructions may cause the machine to fabricate the logic to perform the techniques described herein. Such representations, known as “IP cores,” are reusable units of logic for an integrated circuit that may be stored on a tangible, machine-readable medium as a hardware model that describes the structure of the integrated circuit. The hardware model may be supplied to various customers or manufacturing facilities, which load the hardware model on fabrication machines that manufacture the integrated circuit. The integrated circuit may be fabricated such that the circuit performs operations described in association with any of the embodiments described herein.

[0179] FIG. 11A is a block diagram illustrating an IP core development system 1100 that may be used to manufacture an integrated circuit to perform operations according to an embodiment. The IP core development system 1100 may be used to generate modular, re-usable designs that can be incorporated into a larger design or used to construct an entire integrated circuit (e.g., a system on chip (SOC or SoC) integrated circuit). A design facility 1130 can generate a software simulation 1110 of an IP core design in a high-level programming language (e.g., C/C++). The software simulation 1110 can be used to design, test, and verify the behavior of the IP core using a simulation model 1112. The simulation model 1112 may include functional, behavioral, and/or timing simulations. A register transfer level (RTL) design 1115 can then be created or synthesized from the simulation model 1112. The RTL design 1115 is an abstraction of the behavior of the integrated circuit that models the flow of digital signals between hardware registers, including the associated logic performed using the modeled digital signals. In addition to an RTL design 1115, lower-level designs at the logic level or transistor level may also be created, designed, or synthesized. Thus, the particular details of the initial design and simulation may vary.

[0180] The RTL design 1115 or equivalent may be further synthesized by the design facility into a hardware model 1120, which may be in a hardware description language (HDL), or some other representation of physical design data. The HDL may be further simulated or tested to verify the IP core design. The IP core design can be stored for delivery to a 3rd party fabrication facility 1165 using non-volatile memory 1140 (e.g., hard disk, flash memory, or any non-volatile storage medium). Alternatively, the IP core design may be transmitted (e.g., via the Internet) over a wired connection 1150 or wireless connection 1160. The fabrication facility 1165 may then fabricate an integrated circuit that is based at least in part on the IP core design. The fabricated integrated circuit can be configured to perform operations in accordance with at least one embodiment described herein.

[0181] FIG. 11B illustrates a cross-section side view of an integrated circuit package assembly 1170, according to some embodiments described herein. The integrated circuit package assembly 1170 illustrates an implementation of one or more processor or accelerator devices as described herein. The package assembly 1170 includes multiple units of hardware logic 1172, 1174 connected to a substrate 1180. The logic 1172, 1174 may be implemented at least partly in configurable logic or fixed-functionality logic hardware, and can include one or more portions of any of the processor core(s), graphics processor(s), or other accelerator devices described herein. Each unit of logic 1172, 1174 can be implemented within a semiconductor die and coupled with the substrate 1180 via an interconnect structure 1173. The interconnect structure 1173 may be configured to route electrical signals between the logic 1172, 1174 and the substrate 1180, and can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure 1173 may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic 1172, 1174. In some embodiments, the substrate 1180 is an epoxy-based laminate substrate. The substrate 1180 may include other suitable types of substrates in other embodiments. The package assembly 1170 can be

connected to other electrical devices via a package interconnect 1183. The package interconnect 1183 may be coupled to a surface of the substrate 1180 to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0182] In some embodiments, the units of logic 1172, 1174 are electrically coupled with a bridge 1182 that is configured to route electrical signals between the logic 1172, 1174. The bridge 1182 may be a dense interconnect structure that provides a route for electrical signals. The bridge 1182 may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic 1172, 1174.

[0183] Although two units of logic 1172, 1174 and a bridge 1182 are illustrated, embodiments described herein may include more or fewer logic units on one or more dies. The one or more dies may be connected by zero or more bridges, as the bridge 1182 may be excluded when the logic is included on a single die. Alternatively, multiple dies or units of logic can be connected by one or more bridges. Additionally, multiple logic units, dies, and bridges can be connected together in other possible configurations, including three-dimensional configurations.

[0184] FIG. 11C illustrates a package assembly 1190 that includes multiple units of hardware logic chiplets connected to a substrate 1180. A graphics processing unit, parallel processor, and/or compute accelerator as described herein can be composed from diverse silicon chiplets that are separately manufactured. A diverse set of chiplets with different IP core logic can be assembled into a single device. Additionally, the chiplets can be integrated into a base die or base chiplet using active interposer technology. The concepts described herein enable the interconnection and communication between the different forms of IP within the GPU. IP cores can be manufactured using different process technologies and composed during manufacturing, which avoids the complexity of converging multiple IPs, especially on a large SoC with several flavors IPs, to the same manufacturing process. Enabling the use of multiple process technologies improves the time to market and provides a cost-effective way to create multiple product SKUs. Additionally, the disaggregated IPs are more amenable to being power gated independently, components that are not in use on a given workload can be powered off, reducing overall power consumption.

[0185] In various embodiments a package assembly 1190 can include components and chiplets that are interconnected by a fabric 1185 and/or one or more bridges 1187. The chiplets within the package assembly 1190 may have a 2.5D arrangement using Chip-on-Wafer-on-Substrate stacking in which multiple dies are stacked side-by-side on a silicon interposer 1189 that couples the chiplets with the substrate 1180. The substrate 1180 includes electrical connections to the package interconnect 1183. In one embodiment the silicon interposer 1189 is a passive interposer that includes through-silicon vias (TSVs) to electrically couple chiplets within the package assembly 1190 to the substrate 1180. In one embodiment, silicon interposer 1189 is an active interposer that includes embedded logic in addition to TSVs. In such embodiment, the chiplets within the package assembly 1190 are arranged using 3D face to face die stacking on top of the active interposer 1189. The active interposer 1189 can include hardware logic for I/O 1191, cache memory 1192,

and other hardware logic **1193**, in addition to interconnect fabric **1185** and a silicon bridge **1187**. The fabric **1185** enables communication between the various logic chiplets **1172**, **1174** and the logic **1191**, **1193** within the active interposer **1189**. The fabric **1185** may be a network on chip (NoC) interconnect or another form of packet switched fabric that switches data packets between components of the package assembly. For complex assemblies, the fabric **1185** may be a dedicated chiplet enables communication between the various hardware logic of the package assembly **1190**.

[0186] Bridge structures **1187** within the active interposer **1189** may be used to facilitate a point-to-point interconnect between, for example, logic or I/O chiplets **1174** and memory chiplets **1175**. In some implementations, bridge structures **1187** may also be embedded within the substrate **1180**. The hardware logic chiplets can include special purpose hardware logic chiplets **1172**, logic or I/O chiplets **1174**, and/or memory chiplets **1175**. The hardware logic chiplets **1172** and logic or I/O chiplets **1174** may be implemented at least partly in configurable logic or fixed-functionality logic hardware and can include one or more portions of any of the processor core(s), graphics processor(s), parallel processors, or other accelerator devices described herein. The memory chiplets **1175** can be DRAM (e.g., GDDR, HBM) memory or cache (SRAM) memory. Cache memory **1192** within the active interposer **1189** (or substrate **1180**) can act as a global cache for the package assembly **1190**, part of a distributed global cache, or as a dedicated cache for the fabric **1185**.

[0187] Each chiplet can be fabricated as separate semiconductor die and coupled with a base die that is embedded within or coupled with the substrate **1180**. The coupling with the substrate **1180** can be performed via an interconnect structure **1173**. The interconnect structure **1173** may be configured to route electrical signals between the various chiplets and logic within the substrate **1180**. The interconnect structure **1173** can include interconnects such as, but not limited to bumps or pillars. In some embodiments, the interconnect structure **1173** may be configured to route electrical signals such as, for example, input/output (I/O) signals and/or power or ground signals associated with the operation of the logic, I/O, and memory chiplets. In one embodiment, an additional interconnect structure couples the active interposer **1189** with the substrate **1180**.

[0188] In some embodiments, the substrate **1180** is an epoxy-based laminate substrate. The substrate **1180** may include other suitable types of substrates in other embodiments. The package assembly **1190** can be connected to other electrical devices via a package interconnect **1183**. The package interconnect **1183** may be coupled to a surface of the substrate **1180** to route electrical signals to other electrical devices, such as a motherboard, other chipset, or multi-chip module.

[0189] In some embodiments, a logic or I/O chiplet **1174** and a memory chiplet **1175** can be electrically coupled via a bridge **1187** that is configured to route electrical signals between the logic or I/O chiplet **1174** and a memory chiplet **1175**. The bridge **1187** may be a dense interconnect structure that provides a route for electrical signals. The bridge **1187** may include a bridge substrate composed of glass or a suitable semiconductor material. Electrical routing features can be formed on the bridge substrate to provide a chip-to-chip connection between the logic or I/O chiplet **1174** and a memory chiplet **1175**. The bridge **1187** may also be referred

to as a silicon bridge or an interconnect bridge. For example, the bridge **1187**, in some embodiments, is an Embedded Multi-die Interconnect Bridge (EMIB). In some embodiments, the bridge **1187** may simply be a direct connection from one chiplet to another chiplet.

[0190] FIG. 11D illustrates a package assembly **1194** including interchangeable chiplets **1195**, according to an embodiment. The interchangeable chiplets **1195** can be assembled into standardized slots on one or more base chiplets **1196**, **1198**. The base chiplets **1196**, **1198** can be coupled via a bridge interconnect **1197**, which can be similar to the other bridge interconnects described herein and may be, for example, an EMIB. Memory chiplets can also be connected to logic or I/O chiplets via a bridge interconnect. I/O and logic chiplets can communicate via an interconnect fabric. The base chiplets can each support one or more slots in a standardized format for one of logic or I/O or memory/cache.

[0191] In one embodiment, SRAM and power delivery circuits can be fabricated into one or more of the base chiplets **1196**, **1198**, which can be fabricated using a different process technology relative to the interchangeable chiplets **1195** that are stacked on top of the base chiplets. For example, the base chiplets **1196**, **1198** can be fabricated using a larger process technology, while the interchangeable chiplets can be manufactured using a smaller process technology. One or more of the interchangeable chiplets **1195** may be memory (e.g., DRAM) chiplets. Different memory densities can be selected for the package assembly **1194** based on the power, and/or performance targeted for the product that uses the package assembly **1194**. Additionally, logic chiplets with a different number of type of functional units can be selected at time of assembly based on the power, and/or performance targeted for the product. Additionally, chiplets containing IP logic cores of differing types can be inserted into the interchangeable chiplet slots, enabling hybrid processor designs that can mix and match different technology IP blocks.

Exemplary System on a Chip Integrated Circuit

[0192] FIGS. 12-13B illustrate exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

[0193] FIG. 12 is a block diagram illustrating an exemplary system on a chip integrated circuit **1200** that may be fabricated using one or more IP cores, according to an embodiment. Exemplary integrated circuit **1200** includes one or more application processor(s) **1205** (e.g., CPUs), at least one graphics processor **1210**, and may additionally include an image processor **1215** and/or a video processor **1220**, any of which may be a modular IP core from the same or multiple different design facilities. Integrated circuit **1200** includes peripheral or bus logic including a USB controller **1225**, UART controller **1230**, an SPI/SDIO controller **1235**, and an I²S/I²C controller **1240**. Additionally, the integrated circuit can include a display device **1245** coupled to one or more of a high-definition multimedia interface (HDMI) controller **1250** and a mobile industry processor interface (MIPI) display interface **1255**. Storage may be provided by a flash memory subsystem **1260** including flash memory and

a flash memory controller. Memory interface may be provided via a memory controller **1265** for access to SDRAM or SRAM memory devices. Some integrated circuits additionally include an embedded security engine **1270**.

[0194] FIGS. 13A-13B are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein. FIG. 13A illustrates an exemplary graphics processor **1310** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. FIG. 13B illustrates an additional exemplary graphics processor **1340** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to an embodiment. Graphics processor **1310** of FIG. 13A is an example of a low power graphics processor core. Graphics processor **1340** of FIG. 13B is an example of a higher performance graphics processor core. Each of graphics processor **1310** and graphics processor **1340** can be variants of the graphics processor **1210** of FIG. 12.

[0195] As shown in FIG. 13A, graphics processor **1310** includes a vertex processor **1305** and one or more fragment processor(s) **1315A-1315N** (e.g., **1315A**, **1315B**, **1315C**, **1315D**, through **1315N-1**, and **1315N**). Graphics processor **1310** can execute different shader programs via separate logic, such that the vertex processor **1305** is optimized to execute operations for vertex shader programs, while the one or more fragment processor(s) **1315A-1315N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. The vertex processor **1305** performs the vertex processing stage of the 3D graphics pipeline and generates primitives and vertex data. The fragment processor(s) **1315A-1315N** use the primitive and vertex data generated by the vertex processor **1305** to produce a frame-buffer that is displayed on a display device. In one embodiment, the fragment processor(s) **1315A-1315N** are optimized to execute fragment shader programs as provided for in the OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in the Direct 3D API.

[0196] Graphics processor **1310** additionally includes one or more memory management units (MMUs) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B**. The one or more MMU(s) **1320A-1320B** provide for virtual to physical address mapping for the graphics processor **1310**, including for the vertex processor **1305** and/or fragment processor(s) **1315A-1315N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in the one or more cache(s) **1325A-1325B**. In one embodiment the one or more MMU(s) **1320A-1320B** may be synchronized with other MMUs within the system, including one or more MMUs associated with the one or more application processor(s) **1205**, image processor **1215**, and/or video processor **1220** of FIG. 12, such that each processor **1205-1220** can participate in a shared or unified virtual memory system. The one or more circuit interconnect(s) **1330A-1330B** enable graphics processor **1310** to interface with other IP cores within the SoC, either via an internal bus of the SoC or via a direct connection, according to embodiments.

[0197] As shown FIG. 13B, graphics processor **1340** includes the one or more MMU(s) **1320A-1320B**, cache(s) **1325A-1325B**, and circuit interconnect(s) **1330A-1330B** of the graphics processor **1310** of FIG. 13A. Graphics processor **1340** includes one or more shader core(s) **1355A-1355N**

(e.g., **1355A**, **1355B**, **1355C**, **1355D**, **1355E**, **1355F**, through **1355N-1**, and **1355N**), which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. The unified shader core architecture is also configurable to execute direct compiled high-level GPGPU programs (e.g., CUDA). The exact number of shader cores present can vary among embodiments and implementations. Additionally, graphics processor **1340** includes an inter-core task manager **1345**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **1355A-1355N** and a tiling unit **1358** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

OCCLUSION DETERMINATION

[0198] FIG. 14 depicts an example of a scenario in which a HiZ operation provides an ambiguous result for an occlusion analysis of a source primitive relative to a destination primitive. For a HiZ occlusion test, if the received primitive (source **1400**) minimum z value is more than a destination primitive maximum z value ($\text{src_zmin} > \text{dest_zmax}$), then the HiZ test passes and indicates that source primitive is fully occluded by the destination primitive. For HiZ, if the received primitive (source **1402** or **1404**) minimum z value is less than a destination primitive **1450** maximum z value ($\text{src_zmin} < \text{dest_zmax}$), then the source primitive may be fully occluded by the destination primitive but the HiZ test issues an ambiguous result and a per-pixel depth is performed to determine which pixels of the source primitive are not occluded. However, in the examples of source primitives **1402** and **1404**, the source primitives are entirely occluded and a per-pixel depth test is not needed. Performance of a per-pixel depth test can utilize resources that could be used for other tasks or otherwise be allowed to be in low power mode.

[0199] FIG. 15 depicts an example of a scenario in which a HiZ operation provides an ambiguous result for an occlusion analysis of a source primitive relative to a destination primitive. In this example a minimum source primitive **1500** depth in a dimension (either x or y dimension) is less than a maximum destination primitive **1550** depth in a dimension (either x or y dimension). The HiZ test provides an ambiguous result and a per-pixel depth test is performed comparing the source and destination primitive depths. However, the source primitive **1500** is fully occluded by destination primitive **1550**.

[0200] FIG. 16 depicts an example of a scenario in which a HiZ operation provides an ambiguous result for an occlusion analysis of source primitive **1600** relative to a destination primitive **1650**. In this example a minimum source primitive **1600** depth in a dimension (either x or y) is less than a maximum destination primitive **1650** depth in a dimension (either x or y). However, source primitive **1600** would be visible and not occluded by destination primitive **1650**.

[0201] Various examples can determine whether an incoming primitive (source primitive) can be culled at coarse analysis (e.g., HiZ) based on a slope of the incoming primitive (source primitive) compared to a slope of a nearest visible primitive (destination primitive). The depth (Z) value

for a pixel can be generated by the following depth relationship: $Z = x * Cx + y * Cy + Cref$, where the Cx and Cy denote the slope in the X and Y directions of the primitive respectively. In some examples, if an X direction slope of a source primitive (src_cx) is greater than or equal to a slope of a destination primitive ($dest_cx$) and a Y direction slope of the source primitive (src_cy) is greater than or equal to a slope of the destination primitive ($dest_cy$), the entire source primitive is occluded. In some examples, consideration of a starting reference depth of source and destination primitives in the X and Y directions can be considered, in addition to a comparison of slopes in X and Y directions, to determine whether the entire source primitive is occluded. In other words, where source primitive minimum depth $>$ destination primitive minimum depth, and source maximum depth \geq destination maximum depth and if $src_cx \geq dest_cx$ and $src_cy \geq dest_cy$, then the source primitive plan is occluded.

[0202] Accordingly, at least scenarios in FIGS. 14-16 can be detected as CULL or PASS, and not provide an ambiguous results that would utilize per-pixel depth testing. Points where depth of the source plane is equal to depth of the destination plane (e.g., source plane and destination plane have equal slope and same starting reference depth) indicates a same depth value is going to be maintained and the source plane can be considered occluded. Points where depth of the source plane is equal to depth of the destination plane or deeper than depth of the destination plane (e.g., source plane has greater slope than destination plane and same or deeper starting reference depth) indicates that source plane is deeper than destination plane and the source plane is occluded and the plane equation for the destination plane can be utilized for occlusion analysis or rendering.

[0203] Some examples can perform depth test and fill operation potentially faster with no additional floating point adders or multipliers and can cull the source primitive or the source primitive can pass the depth test.

[0204] FIG. 17 depicts an example system. GPU 1700 can include geometry pipeline 1702, which can perform operations of geometry pipeline 820 described with respect to FIG. 8. For example, geometry pipeline 1702 can perform tessellation of 3D objects based on a coarse geometric model that is provided as input to geometry pipeline 1702. For example, geometry pipeline 1702 can operate at least in accordance with Microsoft® DirectX 11 and DirectX 12, OpenGL, Vulkan, SDL, WebGL, Allegro, or others. Cores 1704 can execute graphics and media operations. In some examples, geometry pipeline 1702 can provide geometric objects to be processed by one or more threads dispatched to cores 1704 to perform image rendering.

[0205] Memory 1710 can include one or more of: one or more registers, one or more cache devices (e.g., level 1 cache (L1), level 2 cache (L2), level 3 cache (L3), last level cache (LLC)), volatile memory device, non-volatile memory device, or persistent memory device. Memory 1710 can include static random access memory (SRAM) memory technology or memory technology consistent with high bandwidth memory (HBM), or double data rate (DDR), among others.

[0206] Primitive data 1712 can include X, Y, and/or Z coordinates in 3D space of corners of primitive (e.g., triangles or rectangles) as well as color data (e.g., RGB), reflectance, texture coordinates, normal vectors, tangent vectors, and other data.

[0207] Depth buffer 1714 can include data generated by a vertex shader executed by geometry pipeline 1702. For example, depth buffer 1714 can include values of pixel depth (z value) from 0 to 1 that indicate a depth of pixels of fragments 1720 and 1722 in X, Y, Z space. Based on pixel depth data for fragment 1720 and fragment 1722 in depth buffer 1714, geometry pipeline 1702 can utilize coarse occlusion circuitry 1703 to determine whether fragment 1720 is fully occluded, not occluded, or occluded in part by fragment 1722. For example, coarse occlusion circuitry 1703 can perform a HiZ analysis to determine whether a fragment 1720 occludes fragment 1722.

[0208] In some examples, based on depth data for fragment 1720 and fragment 1722, geometry pipeline 1702 can determine a slope in the x dimension (Cx), a slope in the y dimension (Cy), and starting point (Cref) for fragments 1720 and 1722. In some examples, set up 829 (FIG. 8) can determine the slope in x dimension (Cx), slope in y dimension (Cy), and Cref for source and destination primitives. For example, geometry pipeline 1702 can perform depth and culling operations and access the slopes in x and y directions and starting reference depths for fragments 1720 and 1722 to perform a HiZ occlusion analysis and perform slope comparison and per-pixel depth comparison.

[0209] Referring again to FIG. 8, geometry pipeline 820 can perform operations in clip/setup 829 to determine primitive starting and ending point in terms of slope (e.g., cx, cy, cref). Rasterizer and depth test component 873 can perform a HiZ operation and store a plane equation for a primitive (e.g., contiguous M×N block of pixels, where M and N are integers). The HiZ occlusion analysis operation can access minimum and maximum depth (Z) values of a primitive. Under HiZ, if the minimum Z value of a source primitive is deeper than the maximum Z value of a destination primitive (e.g., source minimum Z $>$ destination maximum Z), then the destination primitive may entirely occlude the source primitive, so that the source primitive is occluded. Under HiZ, if the destination primitive has a maximum Z value that is deeper than the minimum Z value of a source primitive (e.g., source minimum Z $<$ destination maximum Z), then the source primitive may be partially occluded by the destination primitive and a per-pixel depth comparison to determine which pixels are not occluded.

[0210] Referring again to FIG. 17, various examples of geometry pipeline 1702 can potentially detect whether a source primitive is entirely occluded by a destination primitive such as in the scenarios of FIGS. 14-16 by performing a slope-based occlusion determination in addition or alternative to a HiZ analysis (e.g., $src_min < dest_max$). In some examples, if a slope-based occlusion determination is made that slope of source (src_cx) is greater than or equal to slope of destination ($dest_cx$) and slope of source (src_cy) is greater than or equal to slope of destination ($dest_cy$), and if a starting point depth of the source primitive (src_Cref) is greater than or equal to a starting point depth of destination primitive ($dest_cref$), then depth of the destination primitive (e.g., fragment 1720) will not be greater than depth of the source primitive (e.g., fragment 1722) at any pixel location, so the entire source primitive is occluded. For example, if the slope of the source primitive is less than the slope of the destination primitive but starting and ending depth levels of source primitive are greater than starting and ending depth levels of destination primitive, the depth of the destination primitive (e.g., fragment 1720) will not be greater than depth

of the source primitive (e.g., fragment 1722) at any pixel location, and the entire source primitive is occluded.

[0211] Geometry pipeline 1702 can perform slope-based occlusion determination and HiZ occlusion analysis in parallel. Slope-based occlusion can also be used to conclude source tile completely overdraws destination tile, and write a new tile plane equation for source tile when consider depth test as pass. The slope-based occlusion analysis can determine source completely overwrites the destination plane (source is not occluded), since the pixels, where they are the same depth, can be described with either the destination plane depth equation or source plane depth equation.

[0212] Performing slope-based depth test and HiZ allows cull/pass depth operations at HiZ rate, which can run at 4-8 times the per pixel/sample depth test rate. Hence, scenes with abundant overdraw on a depth buffer, which are bottlenecked by slow occlusion analysis operation, can run significantly faster, and consume less power and bandwidth by not performing per-pixel depth testing. Floating point comparators can be used, which can reduce computational cost of depth testing compared to per-pixel depth testing.

[0213] FIG. 18 depicts an example of scenarios. In scenario 1802, slope-based detection indicates a YES result and indicates that destination primitive fully occludes source primitive. Slope-based detection determines that slopes in x and y axes of the source primitive are greater than slopes in x and y axes of the destination primitive and a starting point of the source primitive is greater than starting point of the destination primitive.

[0214] In scenario 1804, slope-based detection indicates an ambiguous result and a per-pixel depth analysis is performed. Slope-based detection determines that slopes in x and y axes of source primitive are greater than slopes in x and y axes of the destination primitive but a starting point of source primitive (Z2Src) is less than starting point of destination primitive (e.g., nearer).

[0215] In scenario 1806, slope-based detection indicates a YES result and indicates that destination primitive fully occludes source primitive. Slope-based detection determines that slopes in x and y axes of source primitive is less than slopes in x and y axes of the destination primitive but starting and ending points of the source primitive is greater than starting and ending points of destination primitive.

[0216] FIG. 19 depicts other examples of scenarios detected by slope-based occlusion. Occlusion analysis of primitive 1902 can pass based on HiZ because source primitive depth minimum is greater than destination primitive 1950 depth maximum. In this case, source primitive is occluded by destination primitive 1950.

[0217] Occlusion analysis of primitive 1904 can pass based on slope analysis because whereas source primitive depth slope is less than destination primitive 1950 depth slope, starting and ending depths of source primitive are both deeper than those of destination primitive 1950.

[0218] Occlusion analysis of primitive 1906 can pass based on slope analysis, despite minimum depth of source primitive being less than maximum depth of destination primitive, because source primitive depth slope is greater than destination primitive depth slope and starting depth of source primitive is deeper than that of destination primitive 1950.

[0219] Occlusion analysis of source primitive 1908 can be ambiguous based on HiZ because minimum depth of source primitive is less than maximum depth of destination primitive 1950. Moreover, an ambiguous occlusion determination can be made based on slope-based analysis because the slope of source primitive 1908 is less than a slope of destination 1950 and a starting pixel depth of the source primitive 1908 destination is less than a starting pixel depth of destination 1950. A per-pixel depth analysis can be performed to determine which pixels of source primitive 1908 is occluded by destination primitive 1950.

[0220] FIG. 20 depicts an example process that can be performed to determine if a first primitive occludes a second primitive. The process can be performed by a geometry pipeline, raster/depth circuitry, or other circuitry or processor-executed software that perform culling of geometries. At 2002, a determination can be made of slopes of depths and a starting point depth for an MxN tile in a depth buffer of a first primitive (e.g., destination primitive), where M and N are integers. For example, the first primitive can be represented as a three-dimensional plane with a Z plane representing depth. At 2004, a determination can be made of slopes of depths and a starting point depth for an MxN tile in a depth buffer of a second primitive (e.g., source or incoming primitive). Slopes of the first primitive can represent depth value slopes in X and Y directions. For example, the second primitive can be represented as a three-dimensional plane with a Z plane representing depth. Slopes of the second primitive can represent depth value slopes in X and Y directions.

[0221] At 2006, a determination can be made as to whether slopes of the first and second primitives are parallel or the X or Y direction slope of the first primitive is greater than a respective X or Y direction slope of the second primitive. Based on the slopes of the first and second primitives being parallel or the X or Y direction slope of the first primitive is greater than a respective X or Y direction slope of the second primitive, the process can proceed to 2008. Based on the slopes of the first and second primitives are not parallel or the X or Y direction slope of the first primitive is less than a respective X or Y direction slope of the second primitive, the process can proceed to 2012.

[0222] At 2008, a determination can be made if a starting point depth of the second primitive is equal to or deeper than that of the first primitive. Based on the starting point depth of the second primitive being equal to or deeper than that of the first primitive, the process can proceed to 2010, where a determination is made that the first primitive occludes the second primitive. Based on the starting point depth of the second primitive being less than that of the first primitive, the process can proceed to 2014, where a determination is made that a per-pixel depth comparison is to be performed to determine which pixels of the second primitive are occluded or not occluded.

[0223] At 2012, the process can determine if the starting point depth of the second primitive is deeper than that of the first primitive and ending point depth of the second primitive deeper than that of the first primitive. Based on the starting point depth of the second primitive being deeper than that of the first primitive and the ending point depth of the second primitive being deeper than that of the first primitive, the process can proceed to 2010, where a determination is made that the first primitive occludes the second primitive. Based on the starting point depth of the second primitive being less deep than that of the first primitive or the ending point depth of the second primitive being less deep than that of the first primitive, the process can proceed to 2014, where a deter-

mination is made that a per-pixel depth comparison is to be performed to determine which pixels of the second primitive are occluded or not occluded.

[0224] Various examples may be implemented using hardware elements, software elements, or a combination of both. In some examples, hardware elements may include devices, components, processors, microprocessors, circuits, circuit elements (e.g., transistors, resistors, capacitors, inductors, and so forth), integrated circuits, ASICs, PLDs, DSPs, FPGAs, memory units, logic gates, registers, semiconductor device, chips, microchips, chip sets, and so forth. In some examples, software elements may include software components, programs, applications, computer programs, application programs, system programs, machine programs, operating system software, middleware, firmware, software modules, routines, subroutines, functions, methods, procedures, software interfaces, APIs, instruction sets, computing code, computer code, code segments, computer code segments, words, values, symbols, or combination thereof. Determining whether an example is implemented using hardware elements and/or software elements may vary in accordance with any number of factors, such as desired computational rate, power levels, heat tolerances, processing cycle budget, input data rates, output data rates, memory resources, data bus speeds and other design or performance constraints, as desired for a given implementation. A processor can be one or more combination of a hardware state machine, digital control logic, central processing unit, or any hardware, firmware and/or software elements.

[0225] Some examples may be implemented using or as an article of manufacture or at least one computer-readable medium. A computer-readable medium may include a non-transitory storage medium to store logic. In some examples, the non-transitory storage medium may include one or more types of computer-readable storage media capable of storing electronic data, including volatile memory or non-volatile memory, removable or non-removable memory, erasable or non-erasable memory, writeable or re-writeable memory, and so forth. In some examples, the logic may include various software elements, such as software components, programs, applications, computer programs, application programs, system programs, machine programs, operating system software, middleware, firmware, software modules, routines, subroutines, functions, methods, procedures, software interfaces, API, instruction sets, computing code, computer code, code segments, computer code segments, words, values, symbols, or combination thereof.

[0226] According to some examples, a computer-readable medium may include a non-transitory storage medium to store or maintain instructions that when executed by a machine, computing device or system, cause the machine, computing device or system to perform methods and/or operations in accordance with the described examples. The instructions may include any suitable type of code, such as source code, compiled code, interpreted code, executable code, static code, dynamic code, and the like. The instructions may be implemented according to a predefined computer language, manner or syntax, for instructing a machine, computing device or system to perform a certain function. The instructions may be implemented using any suitable high-level, low-level, object-oriented, visual, compiled and/or interpreted programming language.

[0227] One or more aspects of at least one example may be implemented by representative instructions stored on at

least one machine-readable medium which represents various logic within the processor, which when read by a machine, computing device or system causes the machine, computing device or system to fabricate logic to perform the techniques described herein. Such representations, known as “IP cores” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

[0228] The appearances of the phrase “one example” or “an example” are not necessarily all referring to the same example or embodiment. Any aspect described herein can be combined with any other aspect or similar aspect described herein, regardless of whether the aspects are described with respect to the same figure or element. Division, omission or inclusion of block functions depicted in the accompanying figures does not infer that the hardware components, circuits, software and/or elements for implementing these functions would necessarily be divided, omitted, or included in embodiments.

[0229] Some examples may be described using the expression “coupled” and “connected” along with their derivatives. These terms are not necessarily intended as synonyms for each other. For example, descriptions using the terms “connected” and/or “coupled” may indicate that two or more elements are in direct physical or electrical contact with each other. The term “coupled,” however, may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

[0230] The terms “first,” “second,” and the like, herein do not denote any order, quantity, or importance, but rather are used to distinguish one element from another. The terms “a” and “an” herein do not denote a limitation of quantity, but rather denote the presence of at least one of the referenced items. The term “asserted” used herein with reference to a signal denote a state of the signal, in which the signal is active, and which can be achieved by applying any logic level either logic 0 or logic 1 to the signal. The terms “follow” or “after” can refer to immediately following or following after some other event or events. Other sequences of operations may also be performed according to alternative embodiments. Furthermore, additional operations may be added or removed depending on the particular applications. Any combination of changes can be used and one of ordinary skill in the art with the benefit of this disclosure would understand the many variations, modifications, and alternative embodiments thereof.

[0231] Disjunctive language such as the phrase “at least one of X, Y, or Z,” unless specifically stated otherwise, is otherwise understood within the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present. Additionally, conjunctive language such as the phrase “at least one of X, Y, and Z,” unless specifically stated otherwise, should also be understood to mean X, Y, Z, or combination thereof, including “X, Y, and/or Z.”

[0232] Illustrative examples of the devices, systems, and methods disclosed herein are provided below. An embodiment of the devices, systems, and methods may include one or more, and combination of, the examples described below.

[0233] Example 1 includes one or more examples, and includes an apparatus, comprising: a memory to store a minimum depth value of a first primitive and maximum depth value of the first primitive and a minimum depth value of a second primitive and maximum depth value of the second primitive and depth slope data of the first primitive and depth slope data of the second primitive and a graphics processing unit (GPU) to determine whether the second primitive is potentially occluded, fully occluded, or not occluded by the first primitive based on a comparison of the minimum depth value of the first primitive and the maximum depth value of the second primitive and a comparison of the depth slope data of the first primitive and the depth slope data of the second primitive and based on a determination the second primitive is fully occluded by the first primitive, exclude the second primitive from rendering.

[0234] Example 2 includes one or more examples, wherein based on the determination that the second primitive is potentially occluded by the first primitive, the GPU is to perform a per-pixel depth comparison between the first primitive and the second primitive to determine a pixel of the second primitive that is occluded.

[0235] Example 3 includes one or more examples, wherein based on a determination that the second primitive is not occluded by the first primitive, the GPU is to determine whether a third primitive is potentially occluded, fully occluded, or not occluded by the second primitive.

[0236] Example 4 includes one or more examples, wherein: the depth slope data of the first primitive comprises an x-axis slope of depth and a y-axis slope of depth of the first primitive and the depth slope data of the second primitive comprises an x-axis slope of depth and a y-axis slope of depth of the second primitive.

[0237] Example 5 includes one or more examples, wherein: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being greater than the depth slope data of the first primitive and the minimum depth value of the second primitive being equal to or more than the minimum depth value of the first primitive.

[0238] Example 6 includes one or more examples, wherein: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being less than the depth slope data of the first primitive and the minimum depth value of the second primitive being more than the minimum depth of the first primitive and the maximum depth value of the second primitive being more than the maximum depth of the first primitive.

[0239] Example 7 includes one or more examples, wherein: determine that the second primitive is fully occluded by the first primitive is based on the minimum depth value of the second primitive being deeper than the minimum depth value of the first primitive and the maximum depth value of the second primitive being deeper than the maximum depth value of the first primitive.

[0240] Example 8 includes one or more examples, and includes a non-transitory computer-readable medium comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: determine whether a second primitive is potentially occluded, fully occluded, or not occluded by a first primitive based on a comparison of a minimum depth value of the first primitive and a maximum depth value of the second primitive.

tive and a comparison of depth slope data of the first primitive and depth slope data of the second primitive and based on a determination the second primitive is fully occluded by the first primitive, exclude the second primitive from rendering.

[0241] Example 9 includes one or more examples, wherein based on a determination that the second primitive is potentially occluded by the first primitive, the one or more processors are to perform a per-pixel depth comparison between the first primitive and the second primitive to determine a pixel of the second primitive that is occluded.

[0242] Example 10 includes one or more examples, wherein based on a determination that the second primitive is not occluded by the first primitive, the one or more processors are to determine whether a third primitive is potentially occluded, fully occluded, or not occluded by the second primitive.

[0243] Example 11 includes one or more examples, wherein: the depth slope data of the first primitive comprises an x-axis slope of depth and a y-axis slope of depth of the first primitive and the depth slope data of the second primitive comprises an x-axis slope of depth and a y-axis slope of depth of the second primitive.

[0244] Example 12 includes one or more examples, comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being greater than the depth slope data of the first primitive and the minimum depth value of the second primitive being equal to or more than the minimum depth value of the first primitive.

[0245] Example 13 includes one or more examples, wherein: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being less than the depth slope data of the first primitive and the minimum depth value of the second primitive being more than the minimum depth of the first primitive and the maximum depth value of the second primitive being more than the maximum depth of the first primitive.

[0246] Example 14 includes one or more examples, wherein: determine that the second primitive is fully occluded by the first primitive is based on the minimum depth value of the second primitive being deeper than the minimum depth value of the first primitive and the maximum depth value of the second primitive being deeper than the maximum depth value of the first primitive.

[0247] Example 15 includes one or more examples, and includes a method comprising: determining whether a second primitive is potentially occluded, fully occluded, or not occluded by a first primitive based on a comparison of a minimum depth value of the first primitive and a maximum depth value of the second primitive and a comparison of depth slope data of the first primitive and depth slope data of the second primitive and based on a determination the second primitive is fully occluded by the first primitive, excluding the second primitive from rendering.

[0248] Example 16 includes one or more examples, wherein based on a determination that the second primitive is potentially occluded by the first primitive, performing a per-pixel depth comparison between the first primitive and the second primitive to determine a pixel of the second primitive that is occluded.

[0249] Example 17 includes one or more examples, wherein based on a determination that the second primitive is not occluded by the first primitive, determining whether a third primitive is potentially occluded, fully occluded, or not occluded by the second primitive.

[0250] Example 18 includes one or more examples, wherein: the depth slope data of the first primitive comprises an x-axis slope of depth and a y-axis slope of depth of the first primitive and the depth slope data of the second primitive comprises an x-axis slope of depth and a y-axis slope of depth of the second primitive.

[0251] Example 19 includes one or more examples, comprising: determining that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being greater than the depth slope data of the first primitive and the minimum depth value of the second primitive being equal to or more than the minimum depth value of the first primitive.

[0252] Example 20 includes one or more examples, wherein: determining that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being less than the depth slope data of the first primitive and the minimum depth value of the second primitive being more than the minimum depth of the first primitive and the maximum depth value of the second primitive being more than the maximum depth of the first primitive.

1. An apparatus, comprising:

a memory to store a minimum depth value of a first primitive and maximum depth value of the first primitive and a minimum depth value of a second primitive and maximum depth value of the second primitive and depth slope data of the first primitive and depth slope data of the second primitive and

a graphics processing unit (GPU) to determine whether the second primitive is potentially occluded, fully occluded, or not occluded by the first primitive based on a comparison of the minimum depth value of the first primitive and the maximum depth value of the second primitive and a comparison of the depth slope data of the first primitive and the depth slope data of the second primitive and

based on a determination the second primitive is fully occluded by the first primitive, exclude the second primitive from rendering.

2. The apparatus of claim 1, wherein based on the determination that the second primitive is potentially occluded by the first primitive, the GPU is to perform a per-pixel depth comparison between the first primitive and the second primitive to determine a pixel of the second primitive that is occluded.

3. The apparatus of claim 1, wherein based on a determination that the second primitive is not occluded by the first primitive, the GPU is to determine whether a third primitive is potentially occluded, fully occluded, or not occluded by the second primitive.

4. The apparatus of claim 1, wherein:

the depth slope data of the first primitive comprises an x-axis slope of depth and a y-axis slope of depth of the first primitive and

the depth slope data of the second primitive comprises an x-axis slope of depth and a y-axis slope of depth of the second primitive.

5. The apparatus of claim 1, wherein: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being greater than the depth slope data of the first primitive and the minimum depth value of the second primitive being equal to or more than the minimum depth value of the first primitive.

6. The apparatus of claim 1, wherein: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being less than the depth slope data of the first primitive and the minimum depth value of the second primitive being more than the minimum depth of the first primitive and the maximum depth value of the second primitive being more than the maximum depth of the first primitive.

7. The apparatus of claim 1, wherein: determine that the second primitive is fully occluded by the first primitive is based on the minimum depth value of the second primitive being deeper than the minimum depth value of the first primitive and the maximum depth value of the second primitive being deeper than the maximum depth value of the first primitive.

8. A non-transitory computer-readable medium comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to:

determine whether a second primitive is potentially occluded, fully occluded, or not occluded by a first primitive based on a comparison of a minimum depth value of the first primitive and a maximum depth value of the second primitive and a comparison of depth slope data of the first primitive and depth slope data of the second primitive and

based on a determination the second primitive is fully occluded by the first primitive, exclude the second primitive from rendering.

9. The computer-readable medium of claim 8, wherein based on a determination that the second primitive is potentially occluded by the first primitive, the one or more processors are to perform a per-pixel depth comparison between the first primitive and the second primitive to determine a pixel of the second primitive that is occluded.

10. The computer-readable medium of claim 8, wherein based on a determination that the second primitive is not occluded by the first primitive, the one or more processors are to determine whether a third primitive is potentially occluded, fully occluded, or not occluded by the second primitive.

11. The computer-readable medium of claim 8, wherein: the depth slope data of the first primitive comprises an x-axis slope of depth and a y-axis slope of depth of the first primitive and

the depth slope data of the second primitive comprises an x-axis slope of depth and a y-axis slope of depth of the second primitive.

12. The computer-readable medium of claim 8, comprising instructions stored thereon, that if executed by one or more processors, cause the one or more processors to:

determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being greater than the depth slope data of the first primitive and the minimum depth value of the second primitive being equal to or more than the minimum depth value of the first primitive.

- 13.** The computer-readable medium of claim **8**, wherein: determine that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being less than the depth slope data of the first primitive and the minimum depth value of the second primitive being more than the minimum depth of the first primitive and the maximum depth value of the second primitive being more than the maximum depth of the first primitive.
- 14.** The computer-readable medium of claim **8**, wherein: determine that the second primitive is fully occluded by the first primitive is based on the minimum depth value of the second primitive being deeper than the minimum depth value of the first primitive and the maximum depth value of the second primitive being deeper than the maximum depth value of the first primitive.
- 15.** A method comprising:
determining whether a second primitive is potentially occluded, fully occluded, or not occluded by a first primitive based on a comparison of a minimum depth value of the first primitive and a maximum depth value of the second primitive and a comparison of depth slope data of the first primitive and depth slope data of the second primitive and
based on a determination the second primitive is fully occluded by the first primitive, excluding the second primitive from rendering.
- 16.** The method of claim **15**, wherein based on a determination that the second primitive is potentially occluded by the first primitive, performing a per-pixel depth comparison between the first primitive and the second primitive to determine a pixel of the second primitive that is occluded.
- 17.** The method of claim **15**, wherein based on a determination that the second primitive is not occluded by the first primitive, determining whether a third primitive is potentially occluded, fully occluded, or not occluded by the second primitive.
- 18.** The method of claim **15**, wherein:
the depth slope data of the first primitive comprises an x-axis slope of depth and a y-axis slope of depth of the first primitive and
the depth slope data of the second primitive comprises an x-axis slope of depth and a y-axis slope of depth of the second primitive.
- 19.** The method of claim **15**, comprising:
determining that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being greater than the depth slope data of the first primitive and the minimum depth value of the second primitive being equal to or more than the minimum depth value of the first primitive.
- 20.** The method of claim **15**, wherein:
determining that the second primitive is fully occluded by the first primitive based on the depth slope data of the second primitive being less than the depth slope data of the first primitive and the minimum depth value of the second primitive being more than the minimum depth of the first primitive and the maximum depth value of the second primitive being more than the maximum depth of the first primitive.

* * * * *