

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250265806

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Koch; Sebastian et al.

DEVICE AND METHOD FOR GENERATING A GRAPH REPRESENTATION FROM A 3-DIMENSIONAL POINT CLOUD

Abstract

A method for training a first machine learning system for generating a graph representation of objects and their relationships in a 3D environment scene from 3D point cloud input data. For each object and each pair of objects and in the scene initial node feature vectors and initial edge feature vectors are determined from the point cloud input data and are arranged in an initial graph structure. A refined graph structure is determined by a graph neural network. From 2-dimensional image sensor data of the environment scene, feature vectors of the objects are determined by a second machine learning system and feature vectors of the object pairs are determined by a third machine learning system. Parameters of the first machine learning system are adjusted.

Inventors: Koch; Sebastian (Stuttgart, DE), Colosi; Mirco (Stuttgart, DE), Vaskevicius; Narunas (Renningen, DE)

Applicant: Robert Bosch GmbH (Stuttgart, DE)

Family ID: 1000008446345

Appl. No.: 19/036393

Filed: January 24, 2025

Foreign Application Priority Data

EP 24 15 8418.4

Feb. 19, 2024

Publication Classification

Int. Cl.: G06V10/426 (20220101); G06T7/13 (20170101); G06T17/00 (20060101); G06V10/25 (20220101); G06V10/44 (20220101); G06V10/82 (20220101)

U.S. Cl.:

Background/Summary

CROSS REFERENCE

[0001] The present application claims the benefit under 35 U.S.C. § 119 of European Patent Application No. EP 24 15 8418.4 filed on Feb. 19, 2024, which is expressly incorporated herein by reference in its entirety.

FIELD

[0002] The present invention relates to a method of training a machine learning system for generating a graph representation of instances and their relationships in a 3-dimensional environment scene from 3-dimensional point cloud input data, a corresponding computer program and a machine-readable storage medium.

BACKGROUND INFORMATION

[0003] 3-dimensional scene graphs are an emergent graph-based representation facilitating various 3D scene understanding tasks. An advantage of 3-dimensional scene graphs is the ability to represent relationships between scene entities, e.g. instance objects in an indoor scene such as a setting in a closed building or enclosed space such as a room, or an outdoor scene. The exploitation of 3-dimensional scene graphs in e.g., navigation task, decision making or task planning in robotics is, however, limited by their availability.

[0004] Due to their complexity and high-level abstraction, predicting 3-dimensional scene graphs may be hard for learned models relying on a small-scale training data set comprising instances and relationships between the instances. To address this issue, open-vocabulary training settings relying on vision-language models (VLMs) such as CLIP (described in arxiv.org/abs/2103.00020) have been proposed, which may allow handling and generating words and descriptions for objects and their inter-relations that are not restricted to a fixed vocabulary.

[0005] In arxiv.org/abs/2309.15940 an implicit scene graph representation for planning navigation task is proposed, relying on human input in the determination of relationships between instances in the graph. Another approach, described in arxiv.org/abs/2309.16650, for instance relies on 2-dimensional VLMs in order to process input data at inference time.

SUMMARY

[0006] According to a first aspect, the present invention relates to a method of training a first machine learning system for generating a graph representation of instances and their relationships in a 3-dimensional environment scene from 3-dimensional point cloud input data P. A graph representation may be a data structure that represents instances and their relationships in a (3-dimensional) environment scene. The graph representation may comprise nodes and edges, wherein the nodes may refer to or may represent instances. The edges may connect two nodes, respectively, wherein an edge may refer to/represent/indicate a relationship between the two instances associated with the nodes that are connected by the respective edge. An instance may be an object in a (3-dimensional) environment scene and a pair of instances may be a pair of objects in that scene. 3-dimensional point cloud input data may be a collection/group of data points referring to an environment scene, wherein each data point has three spatial coordinates. The first machine learning system comprises two preprocessing networks and a graph neural network. The two preprocessing networks may be given by a PointNet, described in arxiv.org/abs/1612.00593, respectively. Preferably, the two preprocessing networks may be point-encoding networks, which may take 3-dimensional point clouds as input and may determine feature vectors for instances

and/or pairs of instances represented by a (sub) set of points in the 3-dimensional input point cloud. Generally speaking, in the context of the present invention presented herein, a feature vector may be an embedding. An embedding may be a representation of (a subset of) input data, e.g., 3-dimensional point cloud or 2-dimensional image sensor data, in a lower-dimensional space. An embedding may be given by a mathematical transformation that maps input data into a vector space, where each dimension of a vector in said vector space represents a specific feature or characteristic of the input data. According to an example embodiment of the present invention, the method comprises the following steps. In a method step, an initial node feature vector $\phi_{\text{sub}.i}$ is determined from the point cloud input data P by the first preprocessing network for each instance i in the scene. Further, an initial edge feature vector $\phi_{\text{sub}.ij}$ is determined from the point cloud input data by the second preprocessing network for each pair of instances i and j in the scene. An (initial) node feature vector may encode features of an instance and an (initial) edge feature vector may encode features of a pair of instances. In other words, features of a pair of instances may be given by a relationship among or between the respective two instances of the instance pair. A feature of an instance pair i and j may be a spatial, comparative and/or a semantic relationship between these two instances i and j . In a further method step, the initial node and initial edge feature vectors are arranged in an initial graph structure by building triplets $(\phi_{\text{sub}.i}, \phi_{\text{sub}.ij}, \phi_{\text{sub}.j})$. In a following method step a refined graph structure is determined by the graph neural network based on the initial graph structure. The refined graph structure comprises refined node feature vectors and refined edge feature vectors. The graph neural network may process and refine the edge feature vectors as well as the node feature vectors and may project them into a latent space suitable for a training objective defined in the following. In a further method step a feature vector of instance i for each instance i in the scene is determined by a second machine learning system from 2-dimensional image sensor data. The 2-dimensional image sensor data refers to the 3-dimensional environment scene. In other words, the 2-dimensional image sensor data are aligned with the 3-dimensional environment scene. The 2-dimensional image sensor data may hence depict/show the same shapes, poses and/or perspectives of the instances and instance pairs in the scene that may be comprised in the 3-dimensional point cloud input data. The second machine learning system may be a vision-language model (VLM). A VLM may be trained on large datasets containing pairs of different modalities describing the same context, e.g. images and corresponding textual descriptions. During training, a VLM may learn mapping the embeddings of different modalities (such as text and images) into a shared vector space. Similar or related concepts from different modalities may be represented by similar or nearby vectors in the shared space. Furthermore, for each pair of instances i and j a feature vector of instance pair i and j is determined in a further step from the 2-dimensional image sensor data by a third machine learning system. The third machine learning system may be a VLM as well. In a further method step, parameters of the first machine learning system are adjusted with respect to a training objective, wherein the training objective is defined by an optimization of a difference between the refined node feature vector of instance i and the corresponding feature vector of instance i for all instances and/or an optimization of a difference between the refined edge feature vectors of instance i and j and the corresponding feature vectors of instance pair i and j for all instance pairs i and j . Preferably, the training objective comprises minimizing a loss function based on a cosine similarity loss, wherein a cosine similarity between a refined node feature vector of instance i and the corresponding feature vector of instance i for all instances shall be maximized and/or a cosine similarity between a refined edge feature vector of instance i and j and the corresponding feature vector of instance pair i and j for all instance pairs shall be maximized.

[0007] Advantageously, the method of the present invention disclosed herein allows to predict/construct graph representations of instances and their relations in a 3-dimensional environment scene directly from 3-dimensional point cloud input data. These data may be acquired by a LiDAR sensor or a RADAR sensor and may directly, or after a preprocessing to, e.g., remove

artifacts in the raw data, be processed by the first machine learning system to construct a refined graph structure, wherein the refined graph structure may encode instances and relations among the instances in a 3-dimensional environment scene. This allows determining a graph structure when only 3-dimensional point cloud data from a sensor are available. Further, a first machine learning system trained according to a method described herein does not require computationally expensive VLM at inference time to construct 3D maps of a 3-dimensional environment scene from 2-dimensional (video) camera data. In addition, explicit semantic relationships may be predicted as part of the proposed method, wherein storing multiple options per edge describing the relationship may become obsolete.

[0008] According to an example embodiment of the present invention, preferably, the first and the second preprocessing networks may comprise PointNets. A PointNet, arxiv.org/abs/1612.00593, may take a 3-dimensional point cloud as input and determine an embedding based on the input point cloud. In a preprocessing step, point clouds P_i or P_{ij} comprising subgroups of the points of an input point cloud P and referring to instance i or pair of instances i and j , respectively, may be extracted from the input point cloud P by using instance masks $M_{sub.i}$. The preprocessing of the 3-dimensional point cloud P into point clouds $P_{sub.i}$ may be performed with an instance segmentation method, such as Mask3D (arxiv.org/abs/2210.03105). A point cloud $P_{sub.ij}$ may then be determined by selecting all points falling within the union of their respective bounding boxes. A PointNet may then receive point clouds $P_{sub.i}$ and determine corresponding initial node feature vectors from these point clouds. Every point cloud $P_{sub.ij}$ may be concatenated with a mask which may be equal to 1 if the respective point corresponds to instance i , 2 if the point corresponds to instance j and 0 otherwise. The concatenated point cloud $P_{sub.ij}$ may then be received by a PointNet to determine initial edge feature vectors corresponding to and/or encoding relationships between instances.

[0009] Preferably, according to an example embodiment of the present invention the second machine learning system may be an OpenSeg (arxiv.org/abs/2112.12143) model and the third machine learning system may be an InstructBLIP (arxiv.org/abs/2305.06500) model. It may be advantageous to use an OpenSeg model as second machine learning system, because it provides language-aligned feature embeddings for each pixel in the image sensor data. InstructBLIP may be especially suited as third machine learning system due to its aligned image data encoder with its language output decoder. In this way relationships between instance pairs i and j may be encoded and, in consequence, extracted/predicted by the method described herein at a level with good compositional understanding of the scene and especially the relation between the instances of an instance pair i and j .

[0010] Preferably, the 3-dimensional point cloud input data are acquired with a LiDAR sensor, a RADAR sensor, a camera with a depth sensor or video-camera with a depth sensor. In the case where the 3-dimensional point cloud data are acquired with a LiDAR or RADAR sensor, these data may be used directly or after preprocessing for, e.g., removing artifacts in the data, as input data. In case of (video) camera sensor data with depth information, a processing of these images is necessary in order to convert the data to 3-dimensional point cloud input data.

[0011] Preferably, the method may further comprise the step of controlling a robot based on the refined graph structure, wherein the refined node feature vectors and refined edge feature vectors shall be determined by the first machine learning system after adjusting parameters of first machine learning system with respect to the training objective. Examples for a robot comprise industrial robots used in manufacturing, at least partly autonomous vehicles and household or garden robots for cleaning or mowing.

[0012] Preferably, the step of determining by the second machine learning system for each instance i in the scene a feature vector of instance i from the 2-dimensional image sensor data as an input may comprise the following sub-steps for each instance i . In a sub step, a set of k image sensor data comprising instance i may be determined from the 2-dimensional image sensor data. Preferably,

instance i as a whole is visible to a high percentage in each of the k image sensor data.

[0013] Accordingly, the k image sensor data may be chosen as the image sensor data with the highest percentage of visibility of instance i within all image sensor data. In a next sub step, a feature vector of instance i may be determined from each of the k image sensor data comprising instance i . Finally, a feature vector of instance i may be obtained by averaging the k determined feature vectors of instance i .

[0014] Preferably, determining by the third machine learning system for each pair of instances i and j a feature vector of instance pair i and j from the 2-dimensional image sensor data comprises the following sub-steps. In a first step a set of m image sensor data is determined comprising the instance pair i and j from the 2-dimensional image sensor data. The m image sensor data may be determined as the sensor data with the highest visibility of both instances i and j out of all image sensor data. Further, bounding boxes for instance i and instance j are determined in each of the m image sensor data comprising the instance pair i and j . In a next step, each of the m image sensor data with the bounding boxes for instance i and instance j are cropped at n different scales to obtain, for each of the m image sensor data, n different cropped image sensor data. Thereby, each of the cropped image sensor data may comprise the bounding boxes of instance i and instance j . In a next step a feature vector of instance pair i and j is determined from each of the n different cropped image sensor data. Finally, a feature vector of instance pair i and j is obtained by first averaging the n feature vectors from the n different cropped image sensor data for each of the m image sensor data to obtain m averaged feature vectors of instance pair i and j and then averaging the m obtained feature vectors of instance pair i and j .

[0015] Using n different cropped sensor data may be beneficial to encode and extract important context information for a relation between the instances from the sensor data in order to determine a relationship between these two instances. Preferably, the refined node feature vectors and the refined edge feature vectors of the refined graph structure are redetermined after adjusting parameters of the first machine learning system with respect to the training objective. This step may be carried out at inference of the first machine learning system, after training the first machine learning system subject to the training objective according to the previously described method steps. In other words, the trained first machine learning system may at inference time determine refined node feature vectors and refined edge feature vectors from 3-dimensional point cloud input data. Further, a list of candidate instances may be provided, wherein each element of the list of candidate instances is a word or a text describing a possible instance in a 3-dimensional environment scene. A possible instance may not be comprised in an actual 3-dimensional environment scene but may be an instance that typically may occur in 3-dimensional scenes. For instance, in indoor environments, possible instances may be, among others, chair, table, bed, floor, etc., which may be part of the environment in a room or another closed environment. Furthermore, a fourth machine learning system, a third preprocessing network and a fifth machine learning system may be provided. The third preprocessing network may be a QFormer (arxiv.org/abs/2301.12597) network. A QFormer network is well suited to translate the features into token-space of the fifth machine learning system. The fourth machine learning system and the first machine learning system shall map their respective input data to the same embedding space. In other words, the first and the fourth machine learning system shall share the same embedding space. The first machine learning system may receive 3-dimensional point cloud data as an input and may be trained by the training objective to map these input data to a lower dimensional vector representation in a corresponding embedding space, whereas the fourth machine learning system may receive word or text as input and map that input to the same embedding space. Accordingly, also the second machine learning system and the third machine learning system map their respective input data to the same embedding space as the fourth machine learning system. The fourth machine learning system may be a CLIP model. The method may further comprise the following steps. In a step an embedding for each element of the list of candidate instances is

determined by the fourth machine learning system. The list of candidate instances may be determined by a user, wherein the user may input the list of candidate instances at an interface and the input list may be provided as an input to the fourth machine learning system. Alternatively, or in addition, the list of candidate instances or a part of the list of candidate instances may be predefined in terms of a hard-coded list of words or text defining possible instances in an environment scene, wherein the hard-coded list may be provided as an input to the fourth machine learning system. In a further step, a graph structure with labeled nodes may be determined based on the refined graph structure. This may be achieved by assigning for each refined node feature vector with corresponding node in the refined graph structure an element of the candidate list to the corresponding node of the refined node feature vector, wherein the assigned element may be a label of the corresponding node. The assignment may be based on the highest similarity between the refined node feature vector and the embeddings of the elements of the candidate list. In other words, an element of the list may be selected and assigned as a label to a respective node, if its embedding has the highest similarity/smallest difference to the refined node feature vector corresponding to the respective node. As a measure for highest similarity/smallest difference, a cosine similarity scoring between the node feature vectors and the text embeddings from the list may be used. The element of the list with the highest cosine similarity may be the assigned to the node, thereby providing a prediction for the node, wherein the prediction refers to the type or class of instance associated with the node. In a further method step, the third preprocessing network may determine input tokens based on the refined edge feature vectors, a predefined query and relationship prompts. A relationship prompt may comprise the labels that are assigned to the nodes connected by the respective edge in the graph structure with labelled nodes. Based on the determined input tokens, the predefined query and the relationship prompts, the fifth machine learning system may determine in a further step a textual description for each refined edge feature vector. In a next step, a scene graph may be determined from the graph structure with labelled nodes by assigning the determined textual description for each refined edge feature vector to the respective edge of the graph structure with labelled nodes.

[0016] The aforementioned proposed method allows to determine a scene graph with instances in a 3-dimensional scene as nodes and a relationships between an instance pair in the scene as edge between two nodes by a two-step prediction. In a first step, instance text labels, i.e. a textual description of an object in the scene in terms of a word or a few words, may be determined. The textual description, i.e. the word or the few words describing an object/instance in the scene, may also be referred to as instance class. In a second step, inter-instance relations may be predicted/determined by providing the edge feature vector, encoding a relationship between the instances for a given instance pair, and the previously determined instance labels as context for a large language model (LLM). Predicting relationships between instances requires compositional understanding of a scene. This issue is approached by the proposed method by leveraging the generative abilities of a pretrained LLM, which may have a good compositional understanding of an environment scene and may be queried to give a textual description for the relation of instances in the scene.

[0017] Preferably, a relationship between two instances may be a spatial, supportive, semantic and/or comparative relationship.

[0018] Preferably, the method may further comprise the step of validating the scene graph by a user and/or controlling a robot based on the scene graph.

[0019] According to a further aspect, the invention relates to a system configured to carry out the method according to steps and/or features described above.

[0020] According to a further aspect, the invention relates to a computer program with machine-readable instructions, which, when executed on one or several computer(s), cause the computer(s) to perform one of the computer-implemented methods described above and below. Furthermore, according to another aspect, the invention relates to a machine-readable storage medium, on which

the above computer program is stored.

[0021] Example embodiments of the present invention will be discussed with reference to the following figures in more detail.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0022] FIG. 1 shows an exemplary information flowchart for an example embodiment of a method of training a first machine learning system for generating a graph representation of instances and their relationships in a 3-dimensional environment scene, according to the present invention.

[0023] FIG. 2 shows another flowchart for a further example embodiment of a method described herein, according to the present invention.

[0024] FIG. 3 shows another flowchart for a further example embodiment of the present invention.

[0025] FIG. 4 shows another flow chart for a further example embodiment of the present invention.

[0026] FIG. 5 shows another flow chart for a further example embodiment of the present invention.

[0027] FIG. 6 shows a flow chart of a method according to an example embodiment of the present invention.

[0028] FIG. 7 shows a system configured to carry out a method according to example embodiments of the present invention disclosed herein.

DETAILED DESCRIPTION OF EXAMPLE EMBODIMENTS

[0029] FIG. 1 shows a flow chart for an embodiment of a method of training a first machine learning system **10** for generating a graph representation of instances and their relationships in a 3-dimensional environment scene **1** from 3-dimensional point cloud input data **P**. First machine learning system **10** comprises preprocessing network **101** and preprocessing network **102**, which may be given by PointNets. First machine learning system **10** further comprises a graph neural network **103**.

[0030] The first preprocessing network **101** may determine for each instance i in the scene **1** an initial node feature vector d_i from the point cloud input data **P**. Second preprocessing network **102** may determine for each pair of instances i and j in the scene **1** an initial edge feature vector $\phi_{sub.ij}$ from the point cloud input data **P**. Initial node and initial edge feature vectors may be arranged in an initial graph structure by building triplets $(\phi_{sub.1}, \phi_{sub.ij}, \phi_{sub.j})$ and a refined graph structure **11** comprising refined node feature vectors **12** and refined edge feature vectors **13** may be determined by graph neural network **103** based on the initial graph structure. Second machine learning system **20** may determine for each instance i in the scene a feature vector **21** of instance i from 2-dimensional image sensor data **2**. The 2-dimensional image sensor data may be aligned to the 3-dimensional point cloud input data **P** of the environment scene **1**. Third machine learning system **30** may determine for each pair of instances i and j a feature vector **31** of instance pair i and j from the 2-dimensional image sensor data **3**. 2-dimensional image sensor data **2** and **3** may be the same data showing instances in the same pose and the same perspective. However, image sensor data **3** may also be cropped with respect to image sensor data **2**.

[0031] Parameters of the first machine learning system **10** may be adjusted with respect to a training objective, wherein the training objective may be defined by a cosine similarity loss. More generally speaking, the training objective may be defined by an optimization of a difference between the refined node feature vector **12** of instance i and the corresponding feature vector **21** of instance i for all instances and/or an optimization of a difference between refined edge feature vector **13** of instance i and j and the corresponding feature vector **31** of instance pair i and j for all instance pairs.

[0032] FIG. 2 shows a flowchart of an embodiment of a computer-implemented method of training a first machine learning system **10** for generating a graph representation **11** of instances and their

relationships in a 3-dimensional environment scene **1** from 3-dimensional point cloud input data **P**. The 3-dimensional point cloud input data **P** may be provided with class-agnostic instance annotation. In this context, “class-agnostic instance annotation” indicates that the point cloud input data **P** may be grouped by instances in the point cloud but no class-labels are known for the instance-wise grouped points. Accordingly, classes of the instances may not be known at this point, but it may be known which points group together to form an instance. First machine learning system **10** may serve the purpose of determining a graph representation **11** from the input point cloud **P**. First machine learning system **10** may comprise first and second preprocessing networks **101** and **102** and a graph neural network **103**. First and second preprocessing networks **101** and **102** may serve as point encoders that determine feature vectors for instances i and j (case of network **101**) and instance pairs i, j (case of network **102**). First and second preprocessing network **101** and **102** may share the same network architecture. First and second preprocessing network **101** and **102** may be PointNets. However, networks **101** and **102** may share the same architecture, but weights may not be shared. First preprocessing network **101** may determine for each instance i in the scene an initial node feature vector d_i from the point cloud input data **P**. Similarly, second preprocessing network **102** may determine for each pair of instances i and j in the scene an initial edge feature vector $\phi_{sub.ij}$ from the point cloud input data **P**. Determined initial node and initial edge feature vectors may be arranged in an initial graph structure by building triplets $(\phi_{sub.i}, \phi_{sub.ij}, \phi_{sub.j})$ and may be provided as an input to the graph neural network **103**. Graph neural network **103** may determine a refined graph structure comprising refined node feature vectors **12** and refined edge feature vectors **13** based on the initial graph structure. In other words, the initial node and edge feature vectors may be processed and refined by graph neural network **103** and graph neural network **103** may project the initial node and edge feature vectors to refined node **12** and edge **13** feature vectors. Turning to the left hand side of FIG. 2, second machine learning system **20** may determine for each instance i in the scene a feature vector **21** of instance i from 2-dimensional image sensor data **2**. Second machine learning system **20** may be a vision-language model that maps image sensor data into an embedding space. A VLM may map image sensor data as well as textual input data into a joint embedding space, wherein image and text input data showing/describing the same context/sujet are mapped to feature vectors with a high similarity score, i.e., e.g. a high cosine similarity. Preferably, second machine learning system **20** is given by an OpenSeg model. Feature vector **21** of instance i may be also referred to as vision-feature embedding of instance i . Third machine learning system **30** may determine for each pair of instances i and j a feature vector of instance pair i and j from the 2-dimensional image sensor data **3**. Preferably, third machine learning system **30** is given by a vision language model such as, e.g., BLIP/InstructBLIP. Feature vector **31** of instance pair i and j may also be referred to as vision-feature embedding of instance pair i, j . 2-dimensional image sensor data **2** and **3** shall be aligned with the 3-dimensional environment scene, by projecting the 3-dimensional environment scene into the image sensor view of image sensor data system **2** and **3** given their pose in 3D. Image sensor data **2** shall show instances i . Image sensor data **3** shall not only show single instances but pairs of instances i and j . Parameters of first machine learning system **10** may be adjusted with respect to a training objective. The training objective may be defined by an optimization of a difference between the refined node feature vector **12** of instance i and the corresponding feature vector **21** of instance i for all instances, , and/or the training objective may be given by an optimization of a difference between refined edge feature vector **13** of instance pair i and j and the corresponding feature vector **31** of instance pair i and j for all instance pairs, . A cosine similarity loss may be used in the training objective,  and/or , to adjust, i.e. “pull”, the graph feature space/embedding space of the first machine learning system **10** towards the embedding space of the second and third machine learning systems **20** and **30**, i.e. the embedding space of the vision language models. Preferably, second and third machine learning system **20** and **30** shall share the same embedding space.

[0033] FIG. 3 shows an exemplary embodiment of sub-steps in a method described herein, related to the determination of a feature vector **31** of instance pair *i* and *j* from 2-dimensional image sensor data **3** by the third machine learning system **30** for each pair of instances *i* and *j*. In coarse of the determination of feature vector **31**, a set of *m* image sensor data comprising the instance pair *i* and *j* may be determined from the 2-dimensional image sensor data **3**. In a next step, bounding boxes for instance *i* and instance *j* may be determined in each of the *m* image sensor data **3** comprising the instance pair *i* and *j*. In a next step, each of the *m* image sensor data with the bounding boxes for instance *i* and instance *j* may be cropped at *n* different scales to obtain for each of the *m* image sensor data *n* different cropped image sensor data. Each of the *n* cropped image sensor data may thereby comprise the bounding boxes of instance *i* and instance *j*. From each of the *n* different cropped image sensor data, a feature vector $\phi_{sup.n}$ of instance pair *i* and *j* may be determined. By averaging the *n* feature vectors $\phi_{sup.n}$ from the *n* different cropped image sensor data for each of the *m* image sensor data, *m* averaged feature vectors **31a**, **31b**, **31c**, of instance pair *i* and *j* may be obtained. Finally, averaging the *m* obtained feature vectors **31a**, **31b**, **31c**, of instance pair *i* and *j*, feature vector **31** of instance pair *i* and *j* may be obtained.

[0034] FIG. 4 shows specific parts of an embodiment of a method described herein. Refined node feature vectors **12** and the refined edge feature vectors **13** of the refined graph structure **11** shall be refined node feature and refined edge feature vectors re-determined after adjusting parameters of the first machine learning system with respect to the training objective, e.g. according to the embodiment described in context of FIGS. 1-3. List of candidate instances **4** is provided, wherein each element of the list of candidate instances is a word or a text describing a possible instance in a 3-dimensional environment scene. List **4** may, among others, comprise the elements “chair, table, bed, floor, . . .” describing elements possible situated in the 3-dimensional environment scene. Fourth machine learning system **40** may determine an embedding **41** for each element of the list **4** of candidate instances. In other words, the fourth machine leaning system maps the elements of the list **4** to feature vectors/embeddings **41** of a lower dimensional embedding space. Based on the refined graph structure **11**, graph structure with labelled nodes **11A** is determined by assigning for each refined node feature vector **12** with corresponding node in the refined graph structure **11** an element of the candidate list **4** to the corresponding node of the refined node feature vector **12**. In this way, a labelled node **12A**, labeled with the assigned element, may be obtained. The assignment is based on the highest similarity between the refined node feature vector and the embeddings of the elements of the candidate list. In other words, that element of list **4** may selected and assigned to the respective node, whose embedding has the highest similarity/smallest difference to the refined node feature vector **12**. As similarity measure, the cosine similarity may be used. In other words, the list element whose text embedding provided by the fourth machine learning system **40** has the highest cosine similarity with a refined node feature vector **12** may be the prediction/class description for the respective node. Third preprocessing network **60** may determine input tokens for fifth machine learning system **50**, based on refined edge feature vectors **13**, a predefined query **6a** and relationship prompts **6b**. Relationship prompts **6b** may comprise the labels **12A** of the nodes connected by the respective edge in the graph structure with labelled nodes. An exemplary and non-limiting example for a relationship prompt may read: “What is the relationship between [label of node *i*] and [label of node *j*]?”. Third preprocessing network **60** may be a QFormer network. Fifth machine learning system **50** may determine for each refined edge feature vector **13** a textual description **13B** based on the determined input tokens, the predefined query **6a** and the relationship prompts **6b**. Accordingly, scene graph **11B** may be determined from the graph structure with labelled nodes **11A** by assigning the determined textual description **13B** for each refined edge feature vector to the respective edge of the graph structure with labelled nodes. For the method to work out, fourth machine learning system **40** and first machine learning system **50** may map their respective input data to the same embedding space.

[0035] FIG. 5 shows parts of an embodiment of the method described herein. Refined graph

structure **11** with refined node feature vectors and refined edge feature vectors is determined after adjusting parameters of the first machine learning system with respect to the training objective. In other words, refined graph structure **11** is generated by trained machine learning model **10**. [0036] List of candidate instances **4** is provided, wherein each element of the list of candidate instances is a word or a text describing a possible instance in a 3-dimensional environment scene. List **4** may be defined by a user or provided by a system. Fourth machine learning system **40** may determine an embedding **41** for each element of the list **4** of candidate instances. Fourth machine learning system **40** may be given by the language encoding part of a VLM, e.g. of CLIP. Based on the refined graph structure **11** a graph structure **11A** with labelled nodes **12A** is determined by assigning for each refined node feature vector **12A** with corresponding node in the refined graph structure **11** an element of the candidate list **4** to the corresponding node of the refined node feature vector **12** based on the highest similarity between the refined node feature vector **12** and the embeddings of the elements of the candidate list **4**. Graph **11A** may contain textual descriptions/words **12A** describing instances *i* at its nodes. However, graph **11A** still contains the refined edge feature vectors at its edges. Third preprocessing network **60** may determine input tokens for fifth machine learning system **50** based on refined edge feature vectors **13**, a predefined query **6a** and relationship prompts **6b**. Relationship prompt **6b** may comprise the labels of the nodes connected by the respective edge in the graph structure with labelled nodes **11A**. Predefined queries **6a** may be pretrained and may guide the third preprocessing network **60** as well as the fifth machine learning system **50** to attend to relevant parts in the computation. Predefined query **6a** may be given by InstructBLIP pretrained queries. Third preprocessing network **60** may translate the refined edge feature vectors and the relationship prompt into the token space of fifth machine learning system **50**. A non-limiting example for a relationship prompt may be given by “What is the relation between [label of node *i*] and [label of node *j*]?”, wherein the labels are taken from the nodes of graph **11A**. Based on the determined input tokens, the predefined query **6a** and the relationship prompts **6b**, fifth machine learning system may determine for each refined edge feature vector a textual description **13B**. For the method to work out, preferably, first, fourth and fifth machine learning system may map their respective input data to the same embedding space. Fifth machine learning system may be a Vicuna 7B model (lmsys.org/blog/2023-03-30-vicuna) using the Llama architecture (arxiv.org/abs/2302.13971), which may be one of the best open-source language models available. It may be noted that 7B refers to the 7 billion (trained) parameters of the Vicuna model. From the graph structure **11A** with labelled nodes a scene graph **11B** may be determined by assigning the determined textual description **13B** for each refined edge feature vector **13** to the respective edge of the graph structure **12B** with labelled nodes.

[0037] FIG. **6** shows an example embodiment of a method described herein according to the present invention. In method step **S1**, an initial node feature vector $\phi_{\text{sub}.i}$ for each instance *i* in the scene may be determined from the point cloud input data *P* by the first preprocessing network, and an initial edge feature vector $\phi_{\text{sub}.ij}$ for each pair of instances *i* and *j* in the scene may be determined from the point cloud input data by the second preprocessing network. In method step **S2**, initial node and initial edge feature vectors may be arranged in an initial graph structure by building triplets $(\phi_{\text{sub}.i}, \phi_{\text{sub}.ij}, \phi_{\text{sub}.j})$. Method step **S3** may be given by determining a refined graph structure comprising refined node feature vectors and refined edge feature vectors by the graph neural network based on the initial graph structure. In method step **S4a** and **S4b**, which may or may not be carried out in parallel to method steps **S1**, **S2**, and/or **S3**, for each instance *I* in the scene a feature vector of instance *I* may be determined by a second machine learning system from 2-dimensional image sensor data (step **S4a**) and for each pair of instances *i* and *j* a feature vector of instance pair *i* and *j* may be determined by a third machine learning system from the 2-dimensional image sensor data. In method step **S5**, parameters of the first machine learning system may be adjusted with respect to a training objective, wherein the training objective may be defined by an optimization of a difference between the refined node feature vector of instance *i* and the

corresponding feature vector of instance *i* for all instances and/or an optimization of a difference between refined edge feature vector of instance *i* and *j* and the corresponding feature vector of instance pair *i* and *j* for all instance pairs.

[0038] FIG. 7 shows an exemplary embodiment of a data processing system **70**, which comprises at least one processor **71** and at least one machine-readable storage medium **72**, the machine-readable storage medium **72** containing instructions which, when executed by the processor **71**, cause the data processing system **70** to carry out a method according to one of the aspects or embodiments of the invention described herein.

[0039] The term “computer” may be understood as covering any devices for the processing of pre-defined calculation rules. These calculation rules can be in the form of software, hardware or a mixture of software and hardware.

[0040] In general, a plurality can be understood to be indexed, that is, each element of the plurality is assigned a unique index, preferably by assigning consecutive integers to the elements contained in the plurality. Preferably, if a plurality comprises *N* elements, wherein *N* is the number of elements in the plurality, the elements are assigned the integers from 1 to *N*. It may also be understood that elements of the plurality can be accessed by their index.

Claims

1. A computer-implemented method of training a first machine learning system for generating a graph representation of instances and their relationships in a 3-dimensional environment scene from 3-dimensional point cloud input data, wherein the first machine learning system includes two preprocessing networks and a graph neural network, the method comprising the following steps: determining, by a first preprocessing network of the two preprocessor networks, for each instance *i* in the scene an initial node feature vector from the point cloud input data and determining by a second preprocessing network of the two preprocessing networks, for each pair of instances *i* and *j* in the scene, an initial edge feature vector from the point cloud input data; arranging the initial node feature vectors and the initial edge feature vectors in an initial graph structure by building triplets; determining a refined graph structure including refined node feature vectors and refined edge feature vectors by the graph neural network based on the initial graph structure; determining by a second machine learning system, for each instance *i* in the scene, a feature vector of the instance *i* from 2-dimensional image sensor data, wherein the 2-dimensional image sensor data refer to the 3-dimensional environment scene, and determining by a third machine learning system for each pair of instances *i* and *j* a feature vector of the instance pair *i* and *j* from the 2-dimensional image sensor data; adjusting parameters of the first machine learning system with respect to a training objective, wherein the training objective is defined by an optimization of a difference between the refined node feature vector of the instance *i* and the corresponding feature vector of the instance *i* for all instances and/or an optimization of a difference between the refined edge feature vector of the instance pair *i* and *j* and the corresponding feature vector of the instance pair *i* and *j* for all instance pairs.
2. The method according to claim 1, wherein the first and the second preprocessing networks are PointNets.
3. The method according to claim 1, wherein the second machine learning system is an OpenSeg model and wherein the third machine learning system is an InstructBLIP model.
4. The method according to claim 1, wherein the 3-dimensional point cloud input data are acquired with a LiDAR sensor, or a RADAR sensor, or a camera with a depth sensor or a video-camera with a depth sensor.
5. The method according to claim 1, further comprising the following step: controlling a robot based on the refined graph structure, wherein the refined node feature vectors and refined edge feature vectors are determined by the first machine learning system after the adjusting of the

parameters of first machine learning system with respect to the training objective.

6. The method according to claim 1, wherein the determining by the second machine learning system for each instance i in the scene the feature vector of the instance i from the 2-dimensional image sensor data as an input includes the following steps for each instance i : determining a set of k image sensor data including the instance i from the 2-dimensional image sensor data, determining a feature vector of the instance i from each of the k image sensor data including the instance i , obtaining the feature vector of the instance i by averaging the k determined feature vectors of the instance i .

7. The method according to claim 1, wherein the determining by the third machine learning system for each pair of instances i and j the feature vector of the instance pair i and j from the 2-dimensional image sensor data includes the following steps: determining a set of m image sensor data including the instance pair i and j from the 2-dimensional image sensor data, determining bounding boxes for the instance i and the instance j in each of the m image sensor data including the instance pair i and j , cropping each of the m image sensor data with the bounding boxes for the instance i and the instance j at n different scales to obtain for each of the m image sensor data n different cropped image sensor data, wherein each of the cropped image sensor data includes the bounding boxes of the instance i and the instance j , determining a feature vector of the instance pair i and j from each of the n different cropped image sensor data, obtaining the feature vector of the instance pair i and j by first averaging the n feature vectors from the n different cropped image sensor data for each of the m image sensor data to obtain m averaged feature vectors (31a) of instance pair i and j and then averaging the m obtained the feature vectors of instance pair i and j .

8. The method according to claim 1, wherein: the refined node feature vectors and the refined edge feature vectors of the refined graph structure are re-determined after adjusting parameters of the first machine learning system with respect to the training objective, a list of candidate instances is provided, wherein each element of the list of candidate instances is a word or a text describing a possible instance in a 3-dimensional environment scene, and a fourth machine learning system, a third preprocessing network and a fifth machine learning system are provided, wherein the fourth machine learning system and the first machine learning system map their respective input data to the same embedding space; and wherein the method further comprising the following steps: determining by the fourth machine learning system an embedding for each element of the list of candidate instances; determining a graph structure with labelled nodes based on the refined graph structure by assigning for each refined node feature vector with corresponding node in the refined graph structure an element of the candidate list as a label to the corresponding node of the refined node feature vector based on a highest similarity between the refined node feature vector and the embeddings of the elements of the candidate list; determining input tokens by the third preprocessing network based on the refined edge feature vectors, a predefined query, and relationship prompts, wherein each relationship prompt includes the labels of the nodes connected by the respective edge in the graph structure with labelled nodes; determining by the fifth machine learning system for each refined edge feature vector a textual description based on the determined input tokens, the predefined query and the relationship prompts; determining a scene graph from the graph structure with labelled nodes by assigning the determined textual description for each refined edge feature vector to the respective edge of the graph structure with labelled nodes.

9. The method according to claim 8, further comprising the following step: validating the scene graph by a user and/or controlling a robot based on the scene graph.

10. A system configured to train a first machine learning system for generating a graph representation of instances and their relationships in a 3-dimensional environment scene from 3-dimensional point cloud input data, wherein the first machine learning system includes two preprocessing networks and a graph neural network, the system configured to perform the following steps: determining, by a first preprocessing network of the two preprocessor networks, for each instance i in the scene an initial node feature vector from the point cloud input data and

determining by a second preprocessing network of the two preprocessing networks, for each pair of instances i and j in the scene, an initial edge feature vector from the point cloud input data; arranging the initial node feature vectors and the initial edge feature vectors in an initial graph structure by building triplets; determining a refined graph structure including refined node feature vectors and refined edge feature vectors by the graph neural network based on the initial graph structure; determining by a second machine learning system, for each instance i in the scene, a feature vector of the instance i from 2-dimensional image sensor data, wherein the 2-dimensional image sensor data refer to the 3-dimensional environment scene, and determining by a third machine learning system for each pair of instances i and j a feature vector of the instance pair i and j from the 2-dimensional image sensor data; adjusting parameters of the first machine learning system with respect to a training objective, wherein the training objective is defined by an optimization of a difference between the refined node feature vector of the instance i and the corresponding feature vector of the instance i for all instances and/or an optimization of a difference between the refined edge feature vector of the instance pair i and j and the corresponding feature vector of the instance pair i and j for all instance pairs.

11. A non-transitory machine-readable storage medium on which is stored a computer program training a first machine learning system for generating a graph representation of instances and their relationships in a 3-dimensional environment scene from 3-dimensional point cloud input data, wherein the first machine learning system includes two preprocessing networks and a graph neural network, the computer program, when executed by one or more processors, causing the one or more processor to perform the following steps: determining, by a first preprocessing network of the two preprocessor networks, for each instance i in the scene an initial node feature vector from the point cloud input data and determining by a second preprocessing network of the two preprocessing networks, for each pair of instances i and j in the scene, an initial edge feature vector from the point cloud input data; arranging the initial node feature vectors and the initial edge feature vectors in an initial graph structure by building triplets; determining a refined graph structure including refined node feature vectors and refined edge feature vectors by the graph neural network based on the initial graph structure; determining by a second machine learning system, for each instance i in the scene, a feature vector of the instance i from 2-dimensional image sensor data, wherein the 2-dimensional image sensor data refer to the 3-dimensional environment scene, and determining by a third machine learning system for each pair of instances i and j a feature vector of the instance pair i and j from the 2-dimensional image sensor data; adjusting parameters of the first machine learning system with respect to a training objective, wherein the training objective is defined by an optimization of a difference between the refined node feature vector of the instance i and the corresponding feature vector of the instance i for all instances and/or an optimization of a difference between the refined edge feature vector of the instance pair i and j and the corresponding feature vector of the instance pair i and j for all instance pairs.
