US Patent & Trademark Office
Patent Public Search | Text View

United States Patent Application Publication	20250266042
Kind Code	A1
Publication Date	August 21, 2025
Inventor(s)	WEINBERG; Garrett L. et al.

DEVICES AND METHODS FOR INVOKING DIGITAL ASSISTANTS

Abstract

Devices and methods for invoking digital assistants are provided. An example method includes displaying, on a display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a domain-specific digital assistant, and wherein the domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receiving a user voice input; determining whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant: invoking the second digital assistant; determining, based on the user voice input, a first digital assistant response using the second digital assistant; and outputting, via a speaker, the first digital assistant response.

Inventors: WEINBERG; Garrett L. (Santa Cruz, CA), BOOKER; Susan L. (San Carlos, CA), GUZZONI; Didier Rene (Montsur-Rolle, CH), PECK; Jessica (Morgan Hill, CA), SADDLER; Harry J. (Berkeley, CA), STREJA; Lynn Irene (San Francisco, CA), ZHANG; Lily Shuting (Seattle, WA)

Applicant: **Apple Inc.** (Cupertino, CA)

Family ID: 1000008589757

Appl. No.: 19/203991

Filed: May 09, 2025

Related U.S. Application Data

parent US continuation 17682691 20220228 parent-grant-document US 12327554 child US 19203991
us-provisional-application US 63160404 20210312

Publication Classification

Background/Summary

[0001] This application is a continuation of U.S. application Ser. No. 17/682,691, filed on Feb. 28, 2022, entitled "DEVICES AND METHODS FOR INVOKING DIGITAL ASSISTANTS," which claims priority from U.S. Provisional Ser. No. 63/160,404, filed on Mar. 12, 2021, entitled "DEVICES AND METHODS FOR INVOKING DIGITAL ASSISTANTS," which are hereby incorporated by reference in their entirety for all purposes.

FIELD

[0002] This relates generally to digital assistants and, more specifically, to invoking digital assistants based on user input.

BACKGROUND

[0003] Intelligent automated assistants (or digital assistants) can provide a beneficial interface between human users and electronic devices. Such assistants can allow users to interact with devices or systems using natural language in spoken and/or text forms. For example, a user can provide a speech input containing a user request to a digital assistant operating on an electronic device. The digital assistant can interpret the user's intent from the speech input and operationalize the user's intent into tasks. The tasks can then be performed by executing one or more services of the electronic device, and a relevant output responsive to the user request can be returned to the user. In some cases, however, a user may provide a request that either too specific or complex for a system-level digital assistant, particularly when in use with various extended reality technologies.

SUMMARY

[0004] Example methods are disclosed herein. An example method includes displaying, on a display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a domain-specific digital assistant, and wherein the domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receiving a user voice input; determining whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant: invoking the second digital assistant; determining, based on the user voice input, a first digital assistant response using the second digital assistant; and outputting, via a speaker, the first digital assistant response.

[0005] Example non-transitory computer-readable media are disclosed herein. An example non-transitory computer-readable storage medium stores one or more programs. The one or more programs comprise instructions, which when executed by one or more processors of an electronic device having a display, cause the electronic device to display, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a domain-specific digital assistant, and wherein the domain-specific digital assistant corresponds to a first

domain; and while displaying the first avatar within the VR environment: receive a user voice input; determine whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant: invoke the second digital assistant; determine, based on the user voice input, a first digital assistant response using the second digital assistant; and output, via a speaker, the first digital assistant response.

[0006] Example transitory computer-readable media are disclosed herein. An example non-transitory computer-readable storage medium stores one or more programs. The one or more programs comprise instructions, which when executed by one or more processors of an electronic device having a display, cause the electronic device to display, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a domain-specific digital assistant, and wherein the domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receive a user voice input; determine whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant: invoke the second digital assistant; determine, based on the user voice input, a first digital assistant response using the second digital assistant; and output, via a speaker, the first digital assistant response.

[0007] Example electronic devices are disclosed herein. An example electronic device comprises a display, one or more processors; a memory; and one or more programs, where the one or more programs are stored in the memory and configured to be executed by the one or more processors, the one or more programs including instructions for displaying, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a domain-specific digital assistant, and wherein the domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receiving a user voice input; determining whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant: invoking the second digital assistant; determining, based on the user voice input, a first digital assistant response using the second digital assistant; and outputting, via a speaker, the first digital assistant response.

[0008] An example electronic device comprises a display, means for displaying, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a domain-specific digital assistant, and wherein the domain-specific digital assistant corresponds to a first domain; and means for, while displaying the first avatar within the VR environment: receiving a user voice input; determining whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant: invoking the second digital assistant; determining, based on the user voice input, a first digital assistant response using the second digital assistant; and outputting, via a speaker, the first digital assistant response.

[0009] Selectively invoking a domain-specific digital assistant or a second digital assistant allows for a more intuitive and efficient user experience. For example, handling user requests during display of a VR environment allows for requests related to the user's VR experience to be handled without interrupting VR immersion. In turn, this provides for more efficient use of the electronic device (e.g., by avoiding interruption of the display of a VR environment to handle the request), which, additionally, reduces power usage and improves battery life of the device. As another example, handling requests using any of a number of digital assistants (e.g., domain-specific digital

assistants, general digital assistants) avoids processing request with digital assistants that may not be capable of properly handling the requests. In turn, this provides for more efficient use of the electronic device (e.g., reducing number of instances in which requests are improperly handled), which, additionally, reduces power usage and improves battery life of the device.

Description

BRIEF DESCRIPTION OF FIGURES

[0010] FIGS. **1**A-**1**B depict exemplary systems for use in various extended reality technologies, including virtual reality and mixed reality.

[0011] FIG. **2**A is a block diagram illustrating a digital assistant system or a server portion thereof, according to various examples.

[0012] FIG. **2**B illustrates the functions of the digital assistant shown in FIG. **2**A, according to various examples.

[0013] FIG. **2**C illustrates a portion of an ontology, according to various examples.

[0014] FIGS. **3**A-**3**F illustrate an extended reality environment, according to various examples

[0015] FIGS. **4**A-**4**C illustrate an extended reality environment from a perspective of a user, according to various examples.

[0016] FIG. **5** depicts a flow diagram of an exemplary method for invoking a digital assistant, according to various examples.

DESCRIPTION

[0017] Various examples of electronic systems and techniques for using such systems in relation to various extended reality technologies are described.

[0018] Physical settings are those in the world where people can sense and/or interact without use of electronic systems. For example, a room is a physical setting that includes physical elements, such as, physical chairs, physical desks, physical lamps, and so forth. A person can sense and interact with these physical elements of the physical setting through direct touch, taste, sight, smell, and hearing.

[0019] In contrast to a physical setting, an extended reality (XR) setting refers to a computer-produced environment that is partially or entirely generated using computer-produced content. While a person can interact with the XR setting using various electronic systems, this interaction utilizes various electronic sensors to monitor the person's actions, and translates those actions into corresponding actions in the XR setting. For example, if a XR system detects that a person is looking upward, the XR system may change its graphics and audio output to present XR content in a manner consistent with the upward movement. XR settings may respect laws of physics to mimic physical settings.

[0020] Concepts of XR include virtual reality (VR) and augmented reality (AR). Concepts of XR also include mixed reality (MR), which is sometimes used to refer to the spectrum of realities from between physical settings (but not including physical settings) at one end and VR at the other end. Concepts of XR also include augmented virtuality (AV), in which a virtual or computer-produced setting integrates sensory inputs from a physical setting. These inputs may represent characteristics of a physical setting. For example, a virtual object may take on a color captured, using an image sensor, from the physical setting. Or, an AV setting may adopt current weather conditions of the physical setting.

[0021] Some electronic systems for implementing XR operate with an opaque display and one or more imaging sensors for capturing video and/or images of a physical setting. In some implementations, when a system captures images of a physical setting, and displays a representation of the physical setting on an opaque display using the captured images, the displayed images are called a video pass-through. Some electronic systems for implementing XR operate

with a transparent or semi-transparent display (and optionally with one or more imaging sensors). Such a display allows a person to view a physical setting directly through the display, and also allows for virtual content to be added to the person's field of view by superimposing the content and over the physical setting. Some electronic systems for implementing XR operate with a projection system that projects virtual objects onto a physical setting. The projector may present a holograph onto a physical setting, or may project imagery onto a physical surface, or may project onto the eyes (e.g., retina) of a person, for example.

[0022] Electronic systems providing XR settings can have various form factors. A smart phone or tablet computer may incorporate imaging and display components to provide a XR setting. A head mount system may include imaging and display components to provide a XR setting. These systems may provide computing resources for providing XR settings, and may work in conjunction with one another to provide XR settings. For example, a smartphone or a tablet can connect with a head mounted display to provide XR settings. Or, a computer may connect with home entertainment components or vehicular systems to provide an on-window display or a heads-up display. Electronic systems providing XR settings may utilize display technologies such as LEDs, OLEDs, liquid crystal on silicon, a laser scanning light source, a digital light projector, or combinations thereof. Display technologies can employ substrates, through which light is transmitted, including light waveguides, holographic substrates, optical reflectors and combiners, or combinations thereof.

[0023] FIG. **1**A and FIG. **1**B depict exemplary system **100** for use in various extended reality technologies.

[0024] In some examples, as illustrated in FIG. **1**A, system **100** includes device **100***a*. Device **100***a* includes various components, such as processor(s) **102**, RF circuitry(ies) **104**, memory(ies) **106**, image sensor(s) **108**, orientation sensor(s) **110**, microphone(s) **112**, location sensor(s) **116**, speaker(s) **118**, display(s) **120**, and touch-sensitive surface(s) **122**. These components optionally communicate over communication bus(es) **150** of device **100***a*.

[0025] In some examples, elements of system **100** are implemented in a base station device (e.g., a computing device, such as a remote server, mobile device, or laptop) and other elements of system **100** are implemented in a second device (e.g., a head-mounted device). In some examples, device **100***a* is implemented in a base station device or a second device.

[0026] As illustrated in FIG. **1**B, in some examples, system **100** includes two (or more) devices in communication, such as through a wired connection or a wireless connection. First device **100***b* (e.g., a base station device) includes processor(s) **102**, RF circuitry(ies) **104**, and memory(ies) **106**. These components optionally communicate over communication bus(es) **150** of device **100***b*. Second device **100***c* (e.g., a head-mounted device) includes various components, such as processor(s) **102**, RF circuitry(ies) **104**, memory(ies) **106**, image sensor(s) **108**, orientation sensor(s) **110**, microphone(s) **112**, location sensor(s) **116**, speaker(s) **118**, display(s) **120**, and touch-sensitive surface(s) **122**. These components optionally communicate over communication bus(es) **150** of device **100***c*.

[0027] System **100** includes processor(s) **102** and memory(ies) **106**. Processor(s) **102** include one or more general processors, one or more graphics processors, and/or one or more digital signal processors. In some examples, memory(ies) **106** are one or more non-transitory computer-readable storage mediums (e.g., flash memory, random access memory) that store computer-readable instructions configured to be executed by processor(s) **102** to perform the techniques described below.

[0028] System **100** includes RF circuitry(ies) **104**. RF circuitry(ies) **104** optionally include circuitry for communicating with electronic devices, networks, such as the Internet, intranets, and/or a wireless network, such as cellular networks and wireless local area networks (LANs). RF circuitry(ies) **104** optionally includes circuitry for communicating using near-field communication and/or short-range communication, such as Bluetooth®.

[0029] System **100** includes display(s) **120**. Display(s) **120** may have an opaque display. Display(s) **120** may operate with a transparent or semi-transparent display (and optionally with one or more imaging sensors). Such a display allows a person to view a physical setting directly through the display, and also allows for virtual content to be added to the person's field of view by superimposing the content and over the physical setting. Display(s) **120** may incorporate display technologies such as LEDs, OLEDs, liquid crystal on silicon, a laser scanning light source, a digital light projector, or combinations thereof. Display **120** can employ substrates, through which light is transmitted, including light waveguides, holographic substrates, optical reflectors and combiners, or combinations thereof. In one example, the transparent or semi-transparent display may transition selectively between an opaque state and a transparent or semi-transparent state. Other example implementations of display(s) **120** include, but are not limited to heads up displays, display-capable automotive windshields, display-capable windows, display-capable lenses, tablets, smartphones, desktop computers, or laptop computers. Alternatively, system **100** may be designed to interface with an external display (e.g., display of a smartphone). In some examples, system **100** is a projection-based system that uses retinal projection to project images onto the eyes (e.g., retina) of a person or projects virtual objects onto a physical setting (e.g., project imagery onto a physical surface or a holograph onto a physical setting).

[0030] In some examples, system **100** includes touch-sensitive surface(s) **122** for receiving user inputs, such as tap inputs and swipe inputs. In some examples, display(s) **120** and touch-sensitive surface(s) **122** form touch-sensitive display(s).

[0031] System **100** includes image sensor(s) **108**. Image sensors(s) **108** optionally include one or more visible light image sensor, such as charged coupled device (CCD) sensors, and/or complementary metal-oxide-semiconductor (CMOS) sensors operable to obtain images of physical elements from the physical setting. Image sensor(s) also optionally include one or more infrared (IR) sensor(s), such as a passive IR sensor or an active IR sensor, for detecting infrared light from the physical setting. For example, an active IR sensor includes an IR emitter, such as an IR dot emitter, for emitting infrared light into the physical setting. Image sensor(s) **108** also optionally include one or more event camera(s) configured to capture movement of physical elements in the physical setting. Image sensor(s) **108** also optionally include one or more depth sensor(s) configured to detect the distance of physical elements from system **100**. In some examples, system **100** uses CCD sensors, event cameras, and depth sensors in combination to detect the physical setting around system **100**. In some examples, image sensor(s) **108** include a first image sensor and a second image sensor. The first image sensor and the second image sensor are optionally configured to capture images of physical elements in the physical setting from two distinct perspectives. In some examples, system **100** uses image sensor(s) **108** to receive user inputs, such as hand gestures. In some examples, system **100** uses image sensor(s) **108** to detect the position and orientation of system **100** and/or display(s) **120** in the physical setting. For example, system **100** uses image sensor(s) **108** to track the position and orientation of display(s) **120** relative to one or more fixed elements in the physical setting.

[0032] In some examples, system **100** includes microphones(s) **112**. System **100** uses microphone(s) **112** to detect sound from the user and/or the physical setting of the user. In some examples, microphone(s) **112** includes an array of microphones (including a plurality of microphones) that optionally operate in tandem, such as to identify ambient noise or to locate the source of sound in space of the physical setting.

[0033] System **100** includes orientation sensor(s) **110** for detecting orientation and/or movement of system **100** and/or display(s) **120**. For example, system **100** uses orientation sensor(s) **110** to track changes in the position and/or orientation of system **100** and/or display(s) **120**, such as with respect to physical elements in the physical setting. Orientation sensor(s) **110** optionally include one or more gyroscopes and/or one or more accelerometers.

[0034] FIG. **2**A illustrates a block diagram of digital assistant **200** in accordance with various

examples. In some examples, digital assistant **200** is implemented on a standalone computer system. In some examples, digital assistant **200** is distributed across multiple computers and/or systems. In some examples, some of the modules and functions of the digital assistant are divided into a server portion and a client portion, where the client portion resides on one or more user devices (e.g., device **100***a*, **100***b*) or systems (e.g., system **100**) and, optionally, communicates with the server portion through one or more networks, It should be noted that digital assistant **200** is only one example of a digital assistant, and that digital assistant **200** can have more or fewer components than shown, can combine two or more components, or can have a different configuration or arrangement of the components. The various components shown in FIG. **2**A are implemented in hardware, software instructions for execution by one or more processors, firmware, including one or more signal processing and/or application specific integrated circuits, or a combination thereof.

[0035] In some examples, the digital assistant can perform at least some of the following: convert speech input into text; identify a user's intent expressed in a natural language input received from the user; actively elicit and obtain information needed to fully infer the user's intent (e.g., by disambiguating words, games, intentions, etc.); determine the task flow for fulfilling the inferred intent; and execute the task flow to fulfill the inferred intent.

[0036] By way of example, with reference to FIG. **2**B, the digital assistant **200** can obtain user input (e.g., a speech input), for instance, using I/O processing module **228**, and provide responses (e.g., as speech outputs) to the user input. Digital assistant **200** optionally obtains contextual information associated with the user input from the user device, along with or shortly after the receipt of the user input. The contextual information includes user-specific data, vocabulary, and/or preferences relevant to the user input. In some examples, the contextual information also includes software and hardware states of the user device (or system) at the time the user request is received, and/or information related to the surrounding environment of the user at the time that the user request was received. In some examples, the digital assistant **200** also sends follow-up questions to, and receive answers from, the user regarding the user request. When a user request is received by the digital assistant **200** and the user request includes speech input, the digital assistant **200** forwards the speech input to STT processing module **230** (or speech recognizer) for speech-to-text conversions.

[0037] Natural language processing module **232** ("natural language processor") of the digital assistant takes the n-best candidate text representation(s) ("word sequence(s)" or "token sequence(s)") generated by STT processing module **230**, and attempts to associate each of the candidate text representations with one or more "actionable intents" recognized by the digital assistant. An "actionable intent" (or "user intent") represents a task that can be performed by the digital assistant, and can have an associated task flow implemented in task flow models **254**. The associated task flow is a series of programmed actions and steps that the digital assistant takes in order to perform the task. The scope of a digital assistant's capabilities is, in some examples, dependent on the number and variety of task flows that have been implemented and stored in task flow models **254**, or in other words, on the number and variety of "actionable intents" that the digital assistant recognizes. The effectiveness of the digital assistant, however, also dependents on the assistant's ability to infer the correct "actionable intent(s)" from the user request expressed in natural language.

[0038] In some examples, in addition to the sequence of words or tokens obtained from STT processing module **230**, natural language processing module **232** also receives contextual information associated with the user request. The natural language processing module **232** optionally uses the contextual information to clarify, supplement, and/or further define the information contained in the candidate text representations

[0039] In some examples, the natural language processing is based on, e.g., ontology **260**. Ontology **260** is a hierarchical structure containing many nodes, each node representing either an

"actionable intent" or a "property" relevant to one or more of the "actionable intents" or other "properties". A "property" represents a parameter associated with an actionable intent or a sub-aspect of another property. A linkage between an actionable intent node and a property node in ontology **260** defines how a parameter represented by the property node pertains to the task represented by the actionable intent node.

[0040] In some examples, ontology **260** is made up of actionable intent nodes and property nodes. Within ontology **260**, each actionable intent node is linked to one or more property nodes either directly or through one or more intermediate property nodes. Similarly, each property node is linked to one or more actionable intent nodes either directly or through one or more intermediate property nodes. For example, as shown in FIG. **2**C, ontology **260** includes a "restaurant reservation" node (i.e., an actionable intent node). Property nodes "restaurant," "date/time" (for the reservation), and "party size" are each directly linked to the actionable intent node (i.e., the "restaurant reservation" node).

[0041] In addition, property nodes "cuisine," "price range," "phone number," and "location" are sub-nodes of the property node "restaurant," and are each linked to the "restaurant reservation" node (i.e., the actionable intent node) through the intermediate property node "restaurant." For another example, as shown in FIG. **2**C, ontology **260** also includes a "set reminder" node (i.e., another actionable intent node). Property nodes "date/time" (for setting the reminder) and "subject" (for the reminder) are each linked to the "set reminder" node. Since the property "date/time" is relevant to both the task of making a restaurant reservation and the task of setting a reminder, the property node "date/time" is linked to both the "restaurant reservation" node and the "set reminder" node in ontology **260**.

[0042] An actionable intent node, along with its linked property nodes, is described as a "domain." In the present discussion, each domain is associated with a respective actionable intent, and refers to the group of nodes (and the relationships there between) associated with the particular actionable intent. For example, ontology **260** shown in FIG. **2**C includes an example of restaurant reservation domain **262** and an example of reminder domain **264** within ontology **260**. The restaurant reservation domain includes the actionable intent node "restaurant reservation," property nodes "restaurant," "date/time," and "party size," and sub-property nodes "cuisine," "price range," "phone number," and "location." Reminder domain **264** includes the actionable intent node "set reminder," and property nodes "subject" and "date/time." In some examples, ontology **260** is made up of many domains. Each domain shares one or more property nodes with one or more other domains. For example, the "date/time" property node is associated with many different domains (e.g., a scheduling domain, a travel reservation domain, a movie ticket domain, etc.), in addition to restaurant reservation domain **262** and reminder domain **264**.

[0043] While FIG. **2**C illustrates two example domains within ontology **260**, other domains include, for example, "find a movie," "initiate a phone call," "find directions," "schedule a meeting," "send a message," and "provide an answer to a question," "read a list," "providing navigation instructions," "provide instructions for a task" and so on. A "send a message" domain is associated with a "send a message" actionable intent node, and further includes property nodes such as "recipient(s)," "message type," and "message body." The property node "recipient" is further defined, for example, by the sub-property nodes such as "recipient name" and "message address."

[0044] In some examples, ontology **260** includes all domains (and hence actionable intents) that the digital assistant is capable of understanding and acting upon. In some examples, ontology **260** is modified, such as by adding or removing entire domains or nodes, or by modifying relationships between the nodes.

[0045] In some examples, nodes associated with multiple related actionable intents are clustered under a "super domain" in ontology **260**. For example, a "travel" super-domain includes a cluster of property nodes and actionable intent nodes related to travel. In some examples, a digital assistant

may be a domain-specific digital assistant capable of understanding and acting upon a single domain.

[0046] In some examples, each node in ontology **260** is associated with a set of words and/or phrases that are relevant to the property or actionable intent represented by the node. The respective set of words and/or phrases associated with each node are the so-called "vocabulary" associated with the node. The respective set of words and/or phrases associated with each node are stored in vocabulary index **244** in association with the property or actionable intent represented by the node. For example, returning to FIG. **2**B, the vocabulary associated with the node for the property of "restaurant" includes words such as "food," "drinks," "cuisine," "hungry," "eat," "pizza," "fast food," "meal," and so on. For another example, the vocabulary associated with the node for the actionable intent of "initiate a phone call" includes words and phrases such as "call," "phone," "dial," "ring," "call this number," "make a call to," and so on. The vocabulary index **244** optionally includes words and phrases in different languages.

[0047] Natural language processing module **232** receives the candidate text representations (e.g., text string(s) or token sequence(s)) from STT processing module **230**, and for each candidate representation, determines what nodes are implicated by the words in the candidate text representation. In some examples, if a word or phrase in the candidate text representation is found to be associated with one or more nodes in ontology **260** (via vocabulary index **244**), the word or phrase "triggers" or "activates" those nodes. Based on the quantity and/or relative importance of the activated nodes, natural language processing module **232** selects one of the actionable intents as the task that the user intended the digital assistant to perform. In some examples, the domain that has the most "triggered" nodes is selected. In some examples, the domain having the highest confidence value (e.g., based on the relative importance of its various triggered nodes) is selected. In some examples, the domain is selected based on a combination of the number and the importance of the triggered nodes. In some examples, additional factors are considered in selecting the node as well, such as whether the digital assistant has previously correctly interpreted a similar request from a user.

[0048] User data **248** includes user-specific information, such as user-specific vocabulary, user preferences, user address, user's default and secondary languages, user's contact list, and other short-term or long-term information for each user. In some examples, natural language processing module **232** uses the user-specific information to supplement the information contained in the user input to further define the user intent. For example, for a user request "invite my friends to my birthday party," natural language processing module **232** is able to access user data **248** to determine who the "friends" are and when and where the "birthday party" would be held, rather than requiring the user to provide such information explicitly in his/her request.

[0049] In some examples, once natural language processing module **232** identifies an actionable intent (or domain) based on the user request, natural language processing module **232** generates a structured query to represent the identified actionable intent.

[0050] Task flow processing module **236** is configured to receive the structured query (or queries) from natural language processing module **232**, complete the structured query, if necessary, and perform the actions required to "complete" the user's request. In some examples, the various procedures necessary to complete these tasks are provided in task flow models **254**. In some examples, task flow models **254** include procedures for obtaining additional information from the user and task flows for performing actions associated with the actionable intent.

[0051] Once task flow processing module **236** has completed the structured query for an actionable intent, task flow processing module **236** proceeds to perform the task associated with the actionable intent.

[0052] In some examples, natural language processing as described herein may be implemented in a domain-specific manner. By way of example, once a domain has been determined, NLP may be performed by NLP module **232**, or may be performed by one or more domain-specific NLP

modules (which, optionally, may be included in NLP module **232**) to generate a domain-specific structured query (e.g., a structured query that may, optionally, include tasks and/or parameters not included in a corresponding standard structure query for the domain). Additionally or alternatively, such domain-specific structured queries may be handled by either task flow processing module **236** and/or one or more domain-specific task flow processing modules (which, optionally, may be included in task flow processing module **236**).

[0053] Examples described herein are directed to invoking digital assistants. Generally, a user may provide an input (e.g., voice input) including a request that system **100** perform one or more tasks. In response, system **100** may invoke (e.g., selectively invoke) one or more digital assistants, which in turn may process the input (e.g., derive and/or perform the one or more tasks from the input).

[0054] As will be explained in more detail below, in some examples, system **100** may determine whether a user input represents an intent to invoke a digital assistant of system **100**, such as a domain-specific digital assistant associated with an XR environment, another domain-specific digital assistant, or a general digital assistant. In some examples, whether a user input represents an intent to invoke a particular digital assistant of system **100** is determined based on a set of factors, including but not limited to a user input provided by a user, context of the user, context of system **100**, capabilities of one or more digital assistants of system **100**, or any combination thereof. If system **100** determines that a user input represents an intent to invoke a digital assistant, system **100** invokes the digital assistant to process the user input.

[0055] As used herein, a domain-specific digital assistant refers to a digital assistant corresponding to one or more specific domains (e.g., one or more application-specific domains), for instance, of an active ontology (e.g., navigation, weather, restaurant reservations, etc.) such that the domain-specific digital assistant is capable of processing user inputs relating to or associated with the one or more specific domains. In some examples, a domain-specific digital assistant corresponds to a single domain (e.g., travel reservations) and/or one or more subdomains of the domain (e.g., flight reservations, hotel reservations). In some examples, a domain-specific digital assistant corresponds to multiple domains that are, optionally, associated with each other (e.g., movies, TV shows). It will be appreciated that, in some examples, multiple digital assistants may correspond to one or more same domains. For instance, a domain-specific digital assistant and a general digital assistant may both correspond to a weather domain.

[0056] In some examples, a domain-specific digital assistant operates according to a knowledge base that is not available to a general digital assistant (e.g., general digital assistant of system **100**). Accordingly, a domain-specific digital assistant may be capable of processing user inputs (e.g., performing functions and/or actions) that a general digital assistant is not capable of processing, or may be able to more effectively process particular user inputs than a general digital assistant (e.g., process requests more accurately and/or provide more salient results). By way of example, a user input may include a request for weather information. While a general digital assistant of system **100** may be capable of processing the request to provide weather information, a domain-specific digital assistant corresponding to a weather domain may be capable of providing more comprehensive and/or specific weather information than the general digital assistant.

[0057] In some examples, a domain-specific digital assistant is associated with an application (e.g., third party application). Accordingly, domain-specific digital assistants may operate according to a knowledge base provided by the application. As an example, a domain-specific digital assistant associated with a shopping application may be capable of performing more complex tasks for online shopping than a general digital assistant, even if the domain-specific digital assistant and general digital assistant are implemented on a same system or device.

[0058] In some examples, digital assistants may be invoked while an extended reality environment is displayed. FIGS. **3**A-**3**F, for instance, illustrate an exemplary extended reality ("XR") environment **300** in which one or more digital assistants may be invoked, according to various examples. In some examples, XR environment **300** is a VR environment. In other examples, XR

environment **300** is an AR environment. In some examples, XR environment **300** is displayed via a display device (e.g., display(s) **120** of system **100**).

[0059] As shown, XR environment **300** includes avatars **304**, **308**. In some examples, avatar **304** is a representation of the user in the XR environment **300** and/or indicates a location of the user in the XR environment **300**. In some examples, avatar **308** is a representation of digital assistant **309**, and digital assistant **309** may interface with a user in XR environment **300** via avatar **308**. Digital assistant **309** may, for instance, receive inputs from the user, perform tasks derived from the inputs, and provide responses to the user after performing the tasks, as described. While XR environment **300** is shown from a third-person perspective in FIGS. **3**A-**3**F, it will be understood that, in some examples, XR environment **300** is displayed from the perspective of user **304** (FIGS. **4**A-**4**C) such that a user of system **100** can interact with XR environment **300** with a first-person perspective.

[0060] In some examples, XR environment **300** is associated with a software application (e.g., third party software application). XR environment **300** may, for instance, be an XR environment of the application or an XR environment including one or more virtual objects specified by the application. The software application may reside on system **100** or may reside on another device or system remotely accessible by system **100**.

[0061] In some examples, digital assistant **309** is a domain-specific digital assistant associated with an application. Accordingly, digital assistant **309** may correspond to one or more domains of the application and, optionally, operate according to a knowledge base of the application. In this manner, digital assistant **309** may be capable of processing user inputs relating to or associated with the domain(s) of the application, as described. In some examples, XR environment **300** and digital assistant **309** are associated with a same application. As a result, digital assistant **309** may handle application-specific user requests while a user interacts with XR environment **300**.

[0062] FIGS. **3**A-**3**B illustrate an example in which a user input represents an intent to invoke a domain-specific digital assistant of XR environment **300**.

[0063] With reference to FIG. **3**A, in an example operation, system **100** receives user input **306** ("Hey Yogi, what is downward dog?") while displaying the XR environment **300**. User input **306** may be provided as a voice input, but may additionally or alternatively be provided as textual input, gesture, or any combination thereof.

[0064] In some examples, system **100** determines whether a user input represents an intent to invoke a digital assistant. In some examples, determining whether a user input represents an intent to invoke a digital assistant includes determining, for each digital assistant of a set of digital assistants of system **100**, whether the user input represents an intent to invoke the digital assistant. System **100** may, for instance, select a digital assistant of system **100** best suited to process a user input and determine that the user input represents an intent to invoke the selected digital assistant.

[0065] In some examples, system **100** determines a user input represents an intent to invoke a digital assistant if the user input explicitly invokes a particular digital assistant. For example, system **100** may determine whether the user input **306** includes a digital assistant trigger for digital assistant **309** (e.g., "Hey Yogi"). If so, system **100** may determine that the user input represents an intent to invoke digital assistant **309**.

[0066] In some examples, system **100** determines whether a user input represents an intent to invoke a digital assistant based on context of the user. Context of the user may include but is not limited to a direction of a user's gaze, location of the user (e.g., a location of system **100** and/or a location of avatar **304** in XR environment **300**), prior user activity (e.g., user actions performed on system **100**), or any combination thereof. By way of example, system **100** may determine whether a gaze of the user is directed at avatar **308** (e.g., as indicated by the direction of gaze of avatar **304** or by a direction of the user's gaze relative to display **120**), for instance, during a period of time in which the user input **306** is received or within a predetermined period of time after user input **306** is received. If the direction of the gaze of the user satisfies gaze criteria such that system **100** determines that the gaze of the user is directed at avatar **308**, system **100** may determine user input

306 represents an intent to invoke digital assistant **309**.

[0067] In some examples, system **100** determines whether a user input represents an intent to invoke a digital assistant based on capabilities of one or more digital assistants of system **100**. By way of example, system **100** may identify one or more domains corresponding to user input **306** (e.g., using natural-language processing) and determine whether the identified one or more domains match one or more domains of digital assistants of system **100** (e.g., determine if the identified one or more domains include a yoga domain). If, for instance, a domain of user input **306** matches a domain of digital assistant **309**, system **100** may determine user input **306** represents an intent to invoke digital assistant **309**. In some examples, a domain of user input **306** may match a domain of multiple digital assistants. System **100** may identify the digital assistant best suited to process user input **306** and determine that user input **306** represents an intent to invoke the digital assistant determined to be best suited to process user input **306**.

[0068] Additionally or alternatively, system **100** may determine whether the identified one or more domains of a user input match a domain of an application associated with digital assistants of system **100**. In an example in which digital assistant **309** is associated with a yoga application, for instance, system **100** may determine whether the identified one or more domains of user input **306** match a domain of the yoga application (e.g., exercise domain, yoga subdomain, etc.). If system **100** determines a match exists, system **100** may determine that an user input **306** represents an intent to invoke digital assistant **309**.

[0069] With reference to FIG. **3**B, in response to determining that user input **306** represents an intent to invoke digital assistant **309**, system **100** invokes digital assistant **309**. Once invoked, digital assistant **309** processes user input **306** (e.g., performs one or more tasks indicated by a request of user input **306**) and, optionally, generates a response to the request of user input **306**, such as the response **310** ("A downward facing yoga pose").

[0070] In turn, system **100** provides (e.g., outputs) response **310** to the user. With reference to FIG. **4**A, in some examples, response **310** is provided as a visual output (e.g., text displayed in XR environment **300**) and/or as an audible output (e.g., audio provided via a speaker of system **100**). In some examples, response **310** is provided in a manner such that avatar **308** looks to be speaking one or more words of response **310**. System **100** may, for instance, coordinate movement of a mouth of avatar **308** when providing response **310** to simulate avatar **308** speaking response **310**. In some examples, response **310** is provided using a predefined voice setting (e.g., pitch, timbre, volume, accent, tone, etc.) and/or predefined language setting (e.g., US English, UK English, French) corresponding to avatar **308** and/or domain-specific digital assistant **309**. FIGS. **3**C-**3**D illustrate an example in which a user input represents an intent to invoke a domain-specific digital assistant.

[0071] With reference to FIG. **3**C, in an example operation, system **100** receives user input **312** ("Hey FoodBooker, book me a table at Olive Garden") and determines whether user input **312** represents an intent to invoke a digital assistant of system **100**.

[0072] For example, with reference to FIG. **3**D, system **100** may determine that user input **312** includes a digital assistant trigger (e.g., "Hey FoodBooker") for domain-specific digital assistant **315** (e.g., a domain-specific digital assistant for a restaurant reservation domain) and determine that user input **312** represents an intent to invoke domain-specific digital assistant **315**. As another example, system **100** may determine that a domain of user input **312** matches a domain of domain-specific digital assistant **315** and determine that user input **312** represents an intent to invoke domain-specific digital assistant **315**. Both user input **312** and digital assistant **315** may correspond to a restaurant reservation domain, for instance. Additionally or alternatively, system **100** may determine that a domain of user input **312** matches a domain of an application associated with digital assistant **315** and determine that user input **312** represents an intent to invoke domain-specific digital assistant **315**.

[0073] In response to determining that user input **312** represents an intent to invoke domain-

specific digital assistant **315**, system **100** invokes domain-specific digital assistant **315**. With reference to FIG. **4**B, in some examples, invoking digital assistant **315** causes system **100** to display a representation of the digital assistant **315** in XR environment **300**. By way of example, as illustrated in FIG. **3**D, system **100** may display a visual representation of digital assistant **315**, such as avatar **314**, uniquely distinguishing digital assistant **315** from other digital assistants. As shown, in some examples, avatar **308** representing digital assistant **309** and avatar **314** (or another visual representation) representing digital assistant **315** are displayed concurrently. In other examples, system **100** ceases display of avatar **308** when displaying avatar **314** (or another visual representation). In some examples, system **100** forgoes displaying avatar **314** (or another visual representation) when invoking digital assistant **315**.

[0074] In some examples, XR environment **300** is suspended (e.g., held in a static state) during a period of time in which digital assistant **315** is invoked. In this manner, a user may interact with digital assistant **315** without missing content and/or changes to XR environment **300** while interacting with digital assistant **321**. In some examples, suspending XR environment **300** includes suspending the entirety of XR environment **300**. In other examples, suspending XR environment **300** includes suspending a portion of XR environment **300**. System **100** may, for instance, suspend only virtual objects within a threshold distance of a user in some examples.

[0075] Once invoked, digital assistant **315** processes user input **312** and, optionally, generates a response, such as the response **316** ("I've booked you at 7:00 PM for 2 people").

[0076] In turn, system **100** provides (e.g., outputs) response **316** to the user. In some examples, response **316** is provided as a visual output (e.g., text displayed in XR environment **300**) and/or as an audible output (e.g., audio provided via a speaker of system **100**). In some examples, response **316** is provided in a manner such that avatar **314** looks to be speaking one or more words of response **316**. System **100** may, for instance, coordinate movement of a mouth of avatar **314** when providing response **316** to simulate avatar **314** speaking response **316**. In some examples, response **316** is provided using a predefined voice setting (e.g., pitch, timbre, volume, accent, tone, etc.) and/or predefined language setting (e.g., US English, UK English, French) corresponding to either of avatar **314** and domain-specific digital assistant **315**.

[0077] FIGS. **3**E-**3**F illustrate an example in which a user input represents an intent to invoke a general digital assistant.

[0078] With reference to FIG. **3**E, in an example operation, system **100** receives user input **318** ("What time is it"), determines whether the user input **318** represents an intent to invoke a digital assistant of system **100**.

[0079] For example, with reference to FIG. **3**F, system **100** may determine that user input **318** includes a digital assistant trigger for general digital assistant **321** (e.g., "Hey Siri") and determine that user input **318** represents an intent to invoke general digital assistant **321**. As another example, system **100** may determine that a domain of user input **321** does not match a domain of any domain-specific digital assistants of system **100** and determine that user input **318** represents an intent to invoke general digital assistant **321**. As yet another example, system **100** may determine that a domain of user input **318** does not correspond to any software applications of system **100** having a domain-specific digital assistant and determine that user input **318** represents an intent to invoke general digital assistant **321**. As yet another example, system **100** may determine that a domain of user input **318** matches a domain of the active ontology of system **100** and determine that user input **318** represents an intent to invoke general digital assistant **321**.

[0080] In response to determining that user input **318** represents an intent to invoke general digital assistant **321**, system **100** invokes general digital assistant **321**. With reference to FIG. **4**C, in some examples, invoking general digital assistant **321** causes system **100** to display a representation of the digital assistant **321** in XR environment **300**. By way of example, as illustrated in FIG. **3**F, system **100** may display a visual representation of digital assistant **321**, such as avatar **320**, uniquely distinguishing digital assistant **321** from other digital assistants. In other examples, visual

representations of digital assistant **321** include but are not limited to two-dimensional shapes (e.g., circle), three-dimensional shapes (e.g., orb), animations, or any combination thereof. As shown, in some examples, avatar **320** (or another visual representation) representing digital assistant **321** and avatar **308** representing digital assistant **309** are displayed concurrently. In other examples, system **100** ceases display of avatar **308** when displaying avatar **320** (or another visual representation). In some examples, system **100** forgoes displaying avatar **320** (or another visual representation) when invoking digital assistant **321**. In some examples, avatar **320** (or another visual representation) representing digital assistant **321** is always concurrently displayed with avatar **308** in XR environment **300**.

[0081] In some examples, X R environment **300** is suspended (e.g., held in a static state) while digital assistant **321** is invoked. In this manner, a user may interact with digital assistant **321** without missing content and/or changes to XR environment **300** while interacting with digital assistant **321**.

[0082] Once invoked, digital assistant **321** processes user input **318** and, optionally, generates a response, such as the response **322** ("It's 4:15 PM").

[0083] In turn, system **100** provides (e.g., outputs) response **322** to the user. In some examples, response **322** is provided as a visual output (e.g., text displayed in XR environment **300**) and/or as an audible output (e.g., audio provided via a speaker of system **100**). In some examples, response **322** is provided in a manner such that avatar **320** looks to be speaking one or more words of response **322**. System **100** may, for instance, coordinate movement of a mouth of avatar **320** when providing response **322** to simulate avatar **320** speaking response **322**. In some examples, response **322** is provided using a predefined voice setting (e.g., pitch, timbre, volume, accent, tone, etc.) and/or predefined language setting (e.g., US English, UK English, French) corresponding to either of avatar **320** and domain-specific digital assistant **321**.

[0084] While examples set forth herein are generally directed to invoking digital assistants using system **100**, it will be appreciated that digital assistants may be invoked using any system or device(s) including but not limited to mobile phones, smartwatches, desktop computers, laptop computers, tablets, set-top boxes, or any combination thereof. It will further be appreciated that some aspects of the described invention may be implemented using a first device, while other aspects may be implemented using the first device and/or one or more other devices in communication the first device. By way of example, a first device may provide an environment (e.g., XR environment) by which the user interacts with one or more digital assistants, and a second device may perform natural language processing on behalf the one or more digital assistants.

[0085] It will further be appreciated that while described examples are directed to implementing digital assistants in an XR environment, examples of the present invention are not so limited and may be implemented without the use of extended reality. By way of example, system **100** (or another system or device) may select a digital assistant from a set of digital assistants to process a user input when not providing (e.g., displaying) an XR environment.

[0086] FIG. **5** is a flow diagram illustrating a process for invoking a digital assistant, according to various examples. Method **500** is performed at a device (e.g., device **100**). The device can, optionally, include one or more processors, memory, a display, or any combination thereof. In some examples, the device includes one or more input devices (e.g., a touchscreen, a mic, a camera), and a wireless communication radio (e.g., a Bluetooth connection, WiFi connection, a mobile broadband connection, such as a 4G LTE connection). In some examples, the electronic device includes a plurality of cameras. In some embodiments, the electronic device includes only one camera. In some examples, the device includes one or more biometric sensors which, optionally, include a camera, such as an infrared camera, a thermographic camera, or a combination thereof. Some operations in method **500** are, optionally, combined, the orders of some operations are, optionally, changed, and some operations are, optionally, omitted. While the method **500** is described as being implemented with a VR environment, it will be appreciated that the method **500**

may be implemented using any type of XR environment, such as an AR environment.

[0087] At block **502**, an electronic device displays, on a display of the electronic device, a virtual reality ("VR") environment. In some examples, the VR environment is an application-specific VR environment (e.g., a yoga or meditation VR environment for a yoga or meditation software application stored on the electronic device).). In some examples, the VR environment includes a first avatar representing a domain-specific digital assistant. In some examples, a domain-specific digital assistant corresponds to a single software application stored on the electronic device (e.g., the domain-specific digital assistant only responds to user requests or commands relating to or associated with the single software application). In some examples, the domain-specific digital assistant corresponds to a single domain of an active ontology (e.g., navigation, weather, restaurant reservations, etc.) such that the domain-specific digital assistant only responds to user requests or commands relating to or associated with the single domain. In some examples, a domain-specific digital assistant has access to a knowledge base that a general digital assistant of the electronic device does not have access to (e.g., a remotely-stored knowledge base for a third-party software application (e.g., third-party software application data stored on a remote server). In some examples, a domain-specific digital assistant is capable of performing one or more functions and/or actions that a general digital assistant is not capable of performing (e.g., placing an online order for a specific cup of coffee via a Starbucks software application stored on the electronic device). In some examples, the domain-specific digital assistant corresponds to a first domain. In some examples, the domain-specific digital corresponds to a software application, wherein the software application is associated with one or more domains of an active ontology (e.g., a yoga software application is associated with a general exercise/workout domain as well as a more specific yoga domain (e.g., which is a sub-domain of the exercise domain)).

[0088] At block **504**, while displaying the first avatar within the VR environment, the electronic device receives a user voice input (e.g., from a user of the electronic device).

[0089] At block **506**, while displaying the first avatar within the VR environment, the electronic device determines whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant. In some examples, the second digital assistant is a general (e.g., system-wide, domain-agnostic, or the like) digital assistant (e.g., Siri). In some examples, the second digital assistant is a second domain-specific digital assistant. In some examples, this determination includes determining whether the user voice input includes a first digital assistant trigger corresponding to the domain-specific digital assistant or a second digital assistant trigger corresponding to the second digital assistant (e.g., "Hey Siri"). In some examples, this determination includes determining whether a gaze of a user of the electronic device is directed at the first avatar when the user voice input is received or within a predetermined period of time after the user voice input is received. In some examples, this determination includes performing natural language processing based on the user voice input to determine a domain associated with the user voice input, and determining whether the determined domain matches the first domain corresponding to the domain-specific digital assistant or a second domain corresponding to the second digital assistant.

[0090] At block **508**, in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant, the electronic device invokes the second digital assistant digital assistant. In some examples, invoking the second digital assistant includes displaying a digital assistant icon corresponding to the second digital assistant (e.g., a specific icon for the second digital assistant not used by any other digital assistant of the electronic device) within the VR environment. In some examples, invoking the second digital assistant includes displaying a second avatar within the VR environment.

[0091] At block **510**, in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant, the electronic device determines, based on the user voice input, a first digital assistant response using the second digital assistant. (e.g., using an NLP

module of the second digital assistant, a task flow module of the second digital assistant, and/or a services module of the second digital assistant). In some examples, determining the first digital assistant response includes the second digital assistant NLP module (e.g., a general NLP module or a domain-specific NLP module) generating a structured query (e.g., a domain-specific structured query), and the second digital assistant task flow module processing/completing the structured query (and thus performing one or more tasks). In some examples, processing/completing the structured query using the second digital assistant task flow module includes the second digital assistant task flow module working with a second digital assistant services module (e.g., a domain-specific services module) in order to process/complete the structured query).

[0092] At block **512**, in accordance with a determination that the user voice input represents an intent to invoke the second digital assistant, the electronic device outputs, via a speaker, the first digital assistant response (e.g., based on results of one or more tasks performed by the second digital assistant). In some examples, the speaker is a speaker of the electronic device (e.g., a speaker of a mobile phone). In some examples, the speaker is a speaker of a second electronic device that is communicatively connected to the electronic device (e.g., a speaker of a VR headset connected to a mobile device).

[0093] In some examples, determining whether the user voice input represents an intent to invoke the domain-specific digital assistant or an intent to invoke a second digital assistant includes determining whether the user voice input includes a first digital assistant trigger corresponding to the domain-specific digital assistant or a second digital assistant trigger corresponding to the second digital assistant.

[0094] In some examples, in accordance with a determination that the user voice input does not include the first digital assistant trigger or the second digital assistant trigger, the electronic device determines whether a gaze of a user of the electronic device is directed at the first avatar when the user voice input is received (e.g., at a time when the user voice input ends). In some examples, in response to determining that the gaze of the user is directed at the first avatar when the user voice input is received, the electronic device performs natural language processing based on the user voice input to determine a domain associated with the user voice input.

[0095] In some examples, in accordance with a determination that the gaze of the user is not directed at the first avatar when the user voice input is received, the electronic device determines whether the gaze of the user is directed at the first avatar within a predetermined period of time after the user voice input is received (e.g., within 0.5 seconds, 1 second, 2 seconds, or the like). In some examples, in response to determining that the gaze of the user is directed at the first avatar within the predetermined period of time, the electronic device performs natural language processing based on the user voice input to determine a domain associated with the user voice input.

[0096] In some examples, in response to determining that the gaze of the user is not directed at the first avatar within the predetermined period of time, the electronic device forgoes further processing of the user voice input (e.g., forgoing natural language processing, task flow processing, and/or the like).

[0097] In some examples, the electronic device determines whether the determined domain associated with the user voice input matches the first domain corresponding to the domain-specific digital assistant. In some examples, this determination further includes determining whether the determined domain is associated with a software application corresponding to the domain-specific digital assistant. In some examples, a software application is associated with one or more domains of an ontology.

[0098] In some examples, in accordance with a determination that the determined domain matches the first domain, the electronic device determines, based on the user voice input, a second digital assistant response using the domain-specific digital assistant (e.g., using one or more modules of the domain-specific digital assistant that are separate and distinct from modules of the second

digital assistant). In some examples, the electronic device outputs, via the speaker, the second digital assistant response (e.g., instead of the first digital assistant response). In some examples, the electronic device outputs the first audio output such that the first avatar appears to be speaking the first audio output (e.g., the electronic device coordinates movement of a mouth of the first avatar with the output of the first audio output.). In some examples, the electronic device outputs the first audio output based on at least one of a predefined voice setting associated with the domain-specific digital assistant or a predefined language setting associated with the domain-specific digital assistant.

[0099] In some examples, the electronic device invokes the second digital assistant in response to determining that the determined domain does not match the first domain. In some examples, the electronic device invokes the second digital assistant in response to determining that the determined domain is not associated with a software application corresponding to the domain-specific digital assistant (e.g., instead of or in addition to determining that the determined domain does not match the first domain).

[0100] In some examples, in response to determining that the determined domain does not match the first domain and prior to invoking the second digital assistant, the electronic device outputs a first audio output that informs the user of the electronic device that the domain-specific digital assistant is not capable of providing a digital assistant response based on the user voice input. In some examples, the electronic device outputs the first audio output such that the first avatar appears to be speaking the first audio output (e.g., the electronic device coordinates movement of a mouth of the first avatar with the output of the first audio output.). In some examples, the electronic device outputs the first audio output based on at least one of a predefined voice setting associated with the domain-specific digital assistant or a predefined language setting associated with the domain-specific digital assistant. In some examples, in addition to or instead of the first audio output, the electronic device displays, within the VR environment, a first notification (e.g., including text) that informs the user that the domain-specific digital assistant is not capable of providing a digital assistant response based on the user voice input.

[0101] In some examples, the first audio output includes a user query that requests user approval to invoke the second digital assistant. In some examples, the electronic device invokes the second digital assistant in response to receiving a user input (e.g., a second user voice input or a user selection of an affordance displayed in the VR environment) representing a user approval for invoking the second digital assistant.

[0102] In some examples, invoking the second digital assistant includes displaying, within the VR environment, a second avatar representing the second digital assistant (e.g., while still displaying the first avatar). In some examples, invoking the second digital assistant alternatively or additionally includes displaying a digital assistant icon representing the second digital assistant.

[0103] In some examples, the electronic device outputs the first digital assistant response such that the second avatar appears to be speaking the first digital assistant response. In some examples, the electronic device coordinates movement of a mouth of the second avatar with the output of the first digital assistant response.

[0104] In some examples, the electronic device outputs the first digital assistant response based on at least one of a predefined voice setting associated with the second digital assistant and a predefined language setting associated with the second digital assistant. In some examples, each digital assistant may have one or more unique voice characteristics (e.g., accent, pitch, rate, tone, etc.) or language settings (languages, accents, dialects, use of vernacular terms). In some examples, third party applications may specify one or more voice characteristics for digital assistants.

[0105] In some examples, the domain-specific digital assistant determines digital assistant responses only for user inputs (e.g., user voice inputs, user gesture inputs, user text inputs, etc.) associated with the first domain.

[0106] In some examples, the second digital assistant is a system-wide digital assistant that is

capable of determining a digital assistant response for user inputs associated with any domain included in an ontology of the electronic device (e.g., including the first domain).

[0107] In some examples, the second digital assistant is a second domain-specific digital assistant corresponding to a second domain that is different from the first domain, and wherein the second domain-specific digital assistant determines digital assistant responses only for user inputs associated with the second domain.

[0108] Various processes defined herein consider the option of obtaining and utilizing a user's personal information. For example, such personal information may be utilized in order to provide improved digital assistant responses. However, to the extent such personal information is collected, such information should be obtained with the user's informed consent. As described herein, the user should have knowledge of and control over the use of their personal information.

[0109] Personal information will be utilized by appropriate parties only for legitimate and reasonable purposes. Those parties utilizing such information will adhere to privacy policies and practices that are at least in accordance with appropriate laws and regulations. In addition, such policies are to be well-established, user-accessible, and recognized as in compliance with or above governmental/industry standards. Moreover, these parties will not distribute, sell, or otherwise share such information outside of any reasonable and legitimate purposes.

[0110] Users may, however, limit the degree to which such parties may access or otherwise obtain personal information. For instance, settings or other preferences may be adjusted such that users can decide whether their personal information can be accessed by various entities. Furthermore, while some features defined herein are described in the context of using personal information, various aspects of these features can be implemented without the need to use such information. As an example, if user preferences, account names, and/or location history are gathered, this information can be obscured or otherwise generalized such that the information does not identify the respective user.

Claims

1. A non-transitory computer-readable storage medium storing one or more programs, the one or more programs comprising instructions, which when executed by one or more processors of an electronic device having a display, cause the electronic device to: display, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a first domain-specific digital assistant, and wherein the first domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receive a user voice input; identify one or more domains corresponding to the user voice input; determine, based on the identified one or more domains, whether the user voice input represents an intent to invoke the first domain-specific digital assistant or an intent to invoke a second domain-specific digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second domain-specific digital assistant: invoke the second domain-specific digital assistant; determine, based on the user voice input, a first digital assistant response using the second domain-specific digital assistant; and output, via a speaker, the first digital assistant response.

2. The non-transitory computer-readable storage medium of claim 1, wherein determining whether the user voice input represents an intent to invoke the first domain-specific digital assistant or an intent to invoke the second domain-specific digital assistant includes: determining whether the user voice input includes a first digital assistant trigger corresponding to the first domain-specific digital assistant or a second digital assistant trigger corresponding to the second domain-specific digital assistant.

3. The non-transitory computer-readable storage medium of claim 2, wherein the one or more programs further comprise instructions, which when executed by the one or more processors, cause

the electronic device to: in accordance with a determination that the user voice input does not include the first digital assistant trigger or the second digital assistant trigger, determine whether a gaze of a user of the electronic device is directed at the first avatar when the user voice input is received; and in response to determining that the gaze of the user is directed at the first avatar when the user voice input is received, perform natural language processing based on the user voice input to determine a domain associated with the user voice input.

4. The non-transitory computer-readable storage medium of claim 3, wherein the one or more programs further comprise instructions, which when executed by the one or more processors, cause the electronic device to: in accordance with a determination that the gaze of the user is not directed at the first avatar when the user voice input is received, determine whether the gaze of the user is directed at the first avatar within a predetermined period of time after the user voice input is received; and in response to determining that the gaze of the user is directed at the first avatar within the predetermined period of time, perform natural language processing based on the user voice input to determine a domain associated with the user voice input.

5. The non-transitory computer-readable storage medium of claim 4, wherein the one or more programs further comprise instructions, which when executed by the one or more processors, cause the electronic device to: in response to determining that the gaze of the user is not directed at the first avatar within the predetermined period of time, forgo further processing of the user voice input.

6. The non-transitory computer-readable storage medium of claim 3, wherein the one or more programs further comprise instructions, which when executed by the one or more processors, cause the electronic device to: determine whether the determined domain associated with the user voice input matches the first domain corresponding to the first domain-specific digital assistant.

7. The non-transitory computer-readable storage medium of claim 6, wherein the one or more programs further comprise instructions, which when executed by the one or more processors, cause the electronic device to: in accordance with a determination that the determined domain matches the first domain, determine, based on the user voice input, a second digital assistant response using the first domain-specific digital assistant; and output, via the speaker, the second digital assistant response.

8. The non-transitory computer-readable storage medium of claim 6, wherein the electronic device invokes the second digital assistant in response to determining that the determined domain does not match the first domain.

9. The non-transitory computer-readable storage medium of claim 8, wherein the one or more programs further comprise instructions, which when executed by the one or more processors, cause the electronic device to: in response to determining that the determined domain does not match the first domain and prior to invoking the second domain-specific digital assistant: output a first audio output that informs the user of the electronic device that the first domain-specific digital assistant is not capable of providing a digital assistant response based on the user voice input.

10. The non-transitory computer-readable storage medium of claim 9, wherein the first audio output includes a user query that requests user approval to invoke the second domain-specific digital assistant, and wherein the electronic device invokes the second domain-specific digital assistant in response to receiving a user input representing a user approval for invoking the second domain-specific digital assistant.

11. The non-transitory computer-readable storage medium of claim 1, wherein invoking the second domain-specific digital assistant includes displaying, within the VR environment, a second avatar representing the second digital assistant.

12. The non-transitory computer-readable storage medium of claim 11, wherein the electronic device outputs the first digital assistant response such that the second avatar appears to be speaking the first digital assistant response.

13. The non-transitory computer-readable storage medium of claim 1, wherein the electronic device

outputs the first digital assistant response based on at least one of a predefined voice setting associated with the second domain-specific digital assistant and a predefined language setting associated with the second domain-specific digital assistant.

14. The non-transitory computer-readable storage medium of claim 1, wherein the domain-specific digital assistant determines digital assistant responses only for user inputs associated with the first domain.

15. The non-transitory computer-readable storage medium of claim 1, wherein the second domain-specific digital assistant corresponds to a second domain that is different from the first domain, and wherein the second domain-specific digital assistant determines digital assistant responses only for user inputs associated with the second domain.

16. The non-transitory computer-readable storage medium of claim 1, wherein the first domain-specific digital assistant corresponds to a first application and the second domain-specific digital assistant corresponds to a second application, different from the first application.

17. The non-transitory computer-readable storage medium of claim 1, wherein the first domain-specific digital assistant corresponds to a knowledge base unavailable to a general digital assistant.

18. An electronic device comprising: a display; one or more processors; a memory; and one or more programs, wherein the one or more programs are stored in the memory and configured to be executed by the one or more processors, the one or more programs including instructions for: displaying, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a first domain-specific digital assistant, and wherein the first domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receiving a user voice input; identifying one or more domains corresponding to the user voice input; determining, based on the identified one or more domains, whether the user voice input represents an intent to invoke the first domain-specific digital assistant or an intent to invoke a second domain-specific digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second domain-specific digital assistant: invoking the second domain-specific digital assistant; determining, based on the user voice input, a first digital assistant response using the second domain-specific digital assistant; and outputting, via a speaker, the first digital assistant response.

19. A method, comprising: at an electronic device having one or more processors, memory, and a display: displaying, on the display, a virtual reality ("VR") environment, wherein the VR environment includes a first avatar representing a first domain-specific digital assistant, and wherein the first domain-specific digital assistant corresponds to a first domain; and while displaying the first avatar within the VR environment: receiving a user voice input; identifying one or more domains corresponding to the user voice input; determining, based on the identified one or more domains, whether the user voice input represents an intent to invoke the first domain-specific digital assistant or an intent to invoke a second domain-specific digital assistant; and in accordance with a determination that the user voice input represents an intent to invoke the second domain-specific digital assistant: invoking the second domain-specific digital assistant; determining, based on the user voice input, a first digital assistant response using the second domain-specific digital assistant; and outputting, via a speaker, the first digital assistant response.