

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250266045

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

MAIZELS; Aviad et al.

SYNTHESIZING SPEECH FROM FACIAL SKIN MOVEMENTS

Abstract

A method for generating speech includes uploading a reference set of features that were extracted from sensed movements of one or more target regions of skin on faces of one or more reference human subjects in response to words articulated by the subjects and without contacting the one or more target regions. A test set of features is extracted a from the sensed movements of at least one of the target regions of skin on a face of a test subject in response to words articulated silently by the test subject and without contacting the one or more target regions. The extracted test set of features is compared to the reference set of features, and, based on the comparison, a speech output is generated, that includes the articulated words of the test subject.

Inventors: MAIZELS; Aviad (Ramat Hasharon, IL), BARLIYA; Avi (Tel Aviv, IL), WEXLER; Yonatan (Tel Aviv, IL)

Applicant: Q (CUE) LTD. (Ramat Gan, IL)

Family ID: 1000008586757

Assignee: Q (CUE) LTD. (Ramat Gan, IL)

Appl. No.: 19/202227

Filed: May 08, 2025

Related U.S. Application Data

parent US continuation 18180364 20230308 parent-grant-document US 12340808 child US 19202227

parent WO continuation PCT/IB2022/056418 20220712 PENDING child US 18180364

parent US continuation-in-part PCT/US2022/054527 20220516 PENDING child US PCT/IB2022/056418

us-provisional-application US 63229091 20210804

Publication Classification

Int. Cl.: **G10L15/26** (20060101); **G02B27/00** (20060101); **G06F3/01** (20060101); **G06F40/58** (20200101); **G06V10/141** (20220101); **G06V10/60** (20220101); **G06V10/82** (20220101); **G06V40/16** (20220101); **G10L13/02** (20130101); **G10L13/027** (20130101); **G10L13/04** (20130101); **G10L15/16** (20060101); **G10L15/25** (20130101)

U.S. Cl.:

CPC **G10L15/26** (20130101); **G02B27/0093** (20130101); **G06F40/58** (20200101); **G06V10/141** (20220101); **G06V10/60** (20220101); **G06V10/82** (20220101); **G06V40/171** (20220101); **G06V40/174** (20220101); **G06V40/176** (20220101); **G10L13/02** (20130101); **G10L13/027** (20130101); **G10L13/04** (20130101); **G10L15/16** (20130101); **G10L15/25** (20130101); **G06F3/015** (20130101)

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION [0001] This application claims the benefit of U.S. Provisional Patent Application 63/229,091, filed Aug. 4, 2021, and is a continuation in part of PCT Patent Application PCT/IB2022/054527, filed May 16, 2022. The disclosures of both these related applications are incorporated herein by reference.

FIELD OF THE INVENTION

[0002] The present invention relates generally to physiological sensing, and particularly to algorithms, methods and systems for sensing silent human speech.

BACKGROUND

[0003] The process of speech activates nerves and muscles in the chest, neck, and face. Thus, for example, electromyography (EMG) has been used to capture muscle impulses for purposes of silent speech sensing.

SUMMARY

[0004] An embodiment of the present invention that is described hereinafter provides method for generating speech includes uploading a reference set of features that were extracted from sensed movements of one or more target regions of skin on faces of one or more reference human subjects in response to words articulated by the subjects and without contacting the one or more target regions. A test set of features is extracted from the sensed movements of at least one of the target regions of skin on a face of a test subject in response to words articulated silently by the test subject and without contacting the one or more target regions. The extracted test set of features is compared to the reference set of features, and, based on the comparison, a speech output is generated, that includes the articulated words of the test subject.

[0005] In some embodiments, extracting the test features includes extracting the test features without vocalization of the words by the test subject.

[0006] In some embodiments, the test subject and at least one of the reference subjects are the same.

[0007] In an embodiment, extracting the test set of features includes irradiating the one or more target regions of the skin of the test subject with coherent light, and detecting changes in a sensed secondary coherent light pattern due to reflection of the coherent light from the one or more target regions.

[0008] In another embodiment, the uploaded reference set of features and the extracted test set of features each includes a respective waveform calculated for a respective location in a set of locations within the one or more target regions of the skin from a respective time sequence of an energy metric of the sensed secondary coherent light pattern that corresponds to the location.

[0009] In some embodiments, comparing the extracted features includes training and applying a machine learning (ML) algorithm to generate the speech output.

[0010] In some embodiments, generating the speech output includes synthesizing an audio signal corresponding to the speech output.

[0011] In some embodiments, using the speech output, background audio signals are cleaned from a voiced audio signal.

[0012] In an embodiment, generating the speech output includes generating text.

[0013] In another embodiment, generating the speech output includes, upon failing to distinguish in a given time interval between multiple candidate words with at least a predefined confidence level, generating the speech output for the given time interval by mixing audio of two or more of the candidate words.

[0014] In some embodiments, comparing the extracted test set of features to the reference set of features is performed using a trained artificial neural network (ANN), wherein the ANN was trained on a data set collected from a cohort of reference human subjects.

[0015] In some embodiments, the method further includes retraining the ANN using data set collected from test subjects.

[0016] In some embodiments, the method further includes, using the sensed movements of at least one of the target regions of skin on a face of a test subject, indicating an intent of speech by the test subject.

[0017] In an embodiment, the sensed movements are acquired using acquisition rate lower than 200 samples (e.g., frames) per second. In another embodiment, the sensed movements are acquired using acquisition rate between 60 and 140 samples per second.

[0018] In general, acquisition sample rate is lower than 200 samples per second, whatever type of signal is being sampled (e.g., coherent light, microwaves, ultrasound waves, etc.).

[0019] There is additionally provided, in accordance with another embodiment of the present invention, a method for synthesizing speech, the method including receiving input signals from a human subject that are indicative of intended speech by the human subject. The signals are analyzed to extract words corresponding to the intended speech, such that in at least some time intervals of the intended speech, multiple candidate phonemes are extracted together with respective probabilities that each of the candidate phoneme corresponds to the intended speech in a given time interval. Audible speech is synthesized responsively to the extracted phonemes, such that in the at least some of the time intervals, the audible speech is synthesized by mixing the multiple candidate phonemes responsively to the respective probabilities.

[0020] In some embodiments, the input signals include sensed movements of one or more target regions of skin on faces of the human subjects in response to phonemes articulated by the subject and without contacting the one or more target regions.

[0021] In some embodiments, the input signals include at least one of signals received by irradiating the one or more target regions of the skin of the test subject with coherent light, with changes being detected in a sensed secondary coherent light pattern due to reflection of the coherent light from the one or more target regions, one or more of optical lip-readings signals. EMG signals. EEG signals, and noisy audio signals.

[0022] There is further provided, in accordance with another embodiment of the present invention, a system for generating speech, the system including a memory and a processor. The memory is configured to store a reference set of features that were extracted from sensed movements of one or more target regions of skin on faces of one or more reference human subjects in response to words articulated by the subjects and without contacting the one or more target regions. The processor is configured to (i) upload from the memory the reference set of features, (ii) extract a test set of features from the sensed movements of at least one of the target regions of skin on a face of a test subject in response to words articulated silently by the test subject and without contacting the one or more target regions, and (iii) compare the extracted test set of features to the reference set of features, and, based on the comparison, generate a speech output including the articulated words of the test subject.

[0023] In some embodiments, the sensed movements are acquired by optical sensing head and processing circuitry that are fitted inside a stem of wireless headphones.

[0024] There is furthermore provided, in accordance with yet another embodiment of the present

invention, a system for synthesizing speech, the system including a sensor and a processor. The sensor is configured to receive input signals from a human subject that are indicative of intended speech by the human subject. The processor is configured to (a) analyze the signals to extract words corresponding to the intended speech, such that in at least some time intervals of the intended speech, multiple candidate phonemes are extracted by the processor together with respective probabilities that each of the candidate phoneme corresponds to the intended speech in a given time interval, and (b) synthesize audible speech responsively to the extracted phonemes, such that in the at least some of the time intervals, the audible speech is synthesized by the processor by mixing the multiple candidate phonemes responsively to the respective probabilities.

[0025] The present invention will be more fully understood from the following detailed description of the embodiments thereof, taken together with the drawings in which:

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0026] FIG. 1 is a schematic pictorial illustration of a system for silent speech sensing, in accordance with an embodiment of the invention;

[0027] FIG. 2 is a schematic pictorial illustration of a silent speech sensing device, in accordance with another embodiment of the invention;

[0028] FIG. 3 is a block diagram that schematically illustrates functional components of a system for silent speech sensing, in accordance with an embodiment of the invention;

[0029] FIG. 4 is a flow chart that schematically illustrates a method for silent speech sensing, in accordance with an embodiment of the invention;

[0030] FIG. 5 is a flow chart that schematically illustrates a method for training an Artificial Neural Network (ANN) to perform silent speech deciphering, in accordance with an embodiment of the invention;

[0031] FIG. 6 is a flow chart that schematically illustrates a method of using a trained ANN in inference to perform silent speech deciphering, in accordance with an embodiment of the invention;

[0032] FIG. 7 is a flow chart that schematically illustrates a method of preprocessing silent speech sensed data in preparation for speech deciphering, in accordance with an embodiment of the invention; and

[0033] FIG. 8 is a flow chart that schematically illustrates a method of generating an ambiguous speech output for an ambiguous silent speech input, in accordance with an embodiment of the invention.

DETAILED DESCRIPTION OF EMBODIMENTS

[0034] The widespread use of mobile telephones in public spaces creates audio quality issues. For example, when one of the parties in a telephone conversation is in a noisy location, the other party or parties may have difficulty in understanding what they are hearing due to background noise. Moreover, use in public spaces often raises privacy concerns, since conversations are easily overheard by passersby.

[0035] The human brain and neural activity are complex phenomena that involves many human subsystems. One of those subsystems is the facial region which humans use to communicate with others. It is an innate activity that conveys several levels of meaning. At the core, humans communicate with language. The formation of concepts is closely related to the formation of words and then their language-dependent sounds. Humans train to articulate sounds from birth. Even before full language ability evolves, babies use facial expressions, including micro-expressions, to convey deeper information about themselves. The combined interaction with a person provides another value which is trust. While trust in someone begins with their appearance so that we know who we are talking with, their reactions can provide further trust that the person is not incapacitated.

[0036] In the normal process of vocalization, motor neurons activate muscle groups in the face, larynx, and mouth in preparation for propulsion of air flow out of the lungs, and these muscles

continue moving during speech to create words and sentences. Without this air flow, no sounds are emitted from the mouth. Silent speech occurs when the air flow from the lungs is absent, while the muscles in the face, larynx, and mouth continue to articulate the desired sounds. Silent speech can thus be intentional, for example when one articulates words but does not wish to be heard by others. This articulation can occur even when one conceptualizes spoken words without opening our mouths. The resulting activation of our facial muscles gives rise to minute movements of the skin surface. The present disclosure builds on a system for sensing neural activity, the detection focused on the facial region, which allows the readout of residual of muscular activation of the facial region. These muscles are involved in the inter-human communication such as production of sounds, facial expressions, (including micro-expressions). Breathing and other signs humans use for inter-person communication.

[0037] Embodiments of the present invention that are described herein enable users to articulate words and sentences without actually vocalizing the words or uttering any sounds at all. The inventors have found that by properly sensing and decoding these movements, it is possible to reconstruct reliably the actual sequence of words articulated by the user.

[0038] In some embodiments, a system comprising a wearable device and dedicated software tools deciphers data sensed from fine movements of the skin and subcutaneous nerves and muscles on a subject's face, occurring in response to words articulated by the subject with or without vocalization, and use the deciphered words in generating a speech output including the articulated words. Details of devices and methods used in sensing the data from fine movements of the skin, are described in the above-mentioned International Patent Application PCT/IB2022/054527.

[0039] The disclosed deciphering techniques enable users to communicate with others or to record their own thoughts silently, in a manner that is substantially imperceptible to other parties and is also insensitive to ambient noise.

[0040] Some embodiments use sensing devices having the aforementioned form of common consumer wearable items, such as a clip-on headphone or spectacles. In these embodiments, an optical sensing head is held in a location in proximity to the user's face by a bracket that fits in or over the user's ear. The optical sensing head senses coherent light reflected from the face, for example by directing coherent light toward an area of the face, such as the cheek, and sensing changes in the coherent light pattern that arises due to reflection of the coherent light from the face. Processing circuitry in the device processes the signal output by the optical sensing head due to the reflected light to generate a corresponding speech output. In one embodiment, the optical sensing head and processing circuitry is fitted inside a stem of wireless headphones, such as AirPods. In that embodiment, sensing is slightly further away from the sensed skin location and the sensor's viewing angle is typically narrow.

[0041] Alternatively, the disclosed deciphering technique can be used with a silent speech sensing module, including a coherent light source and sensors, may be integrated into a mobile communication device, such as a smartphone. This integrated sensing module senses silent speech when the user holds the mobile communication device in a suitable location in proximity to the user's face.

[0042] In one example, deciphering of silent speech is performed using a machine learning (ML) algorithm, such as a trained artificial neural network (ANN). In this example image processing software converts acquired signals into preprocessed signals, and the trained ANN specifies speech words contained in the preprocessed signals. Different types ANNs may be used, such as a classification NN that eventually outputs words, and a sequence-to-sequence NN which outputs a sentence (word sequence). To train the ANNs, at least several thousands of examples should typically be gathered and augmented, as described above. This "global" training, that relies on a large group of persons (e.g., a cohort of reference human subjects), allows later for a device of a specific user to perform fine adjustments of its deciphering software. In this manner, within minutes or less of wearing the device and turning on the application, the system (e.g., mobile phone and the wearable device) is ready for deciphering.

[0043] In many cases, speech recognition algorithms output some ambiguous results, as described below. In the case of Human-Human real time communication implantation, waiting for the sentence to complete before synthesizing the text to speech to mitigate the ambiguity will result with a significant delay that might not be acceptable. To solve this issue, the disclosed speech synthesizer is configured to quickly generate an ambiguous output for an ambiguous input, so as not to disrupt the natural flow of conversation. The ambiguity itself may still be resolved at a later stage. In some examples, a processor is used for synthesizing speech, by performing the steps of (i) receiving input signals from a human subject that are indicative of intended speech by the human subject, (ii) analyzing the signals to extract words corresponding to the intended speech, such that in at least some time intervals of the intended speech, multiple candidate phonemes are extracted together with respective probabilities that each of the candidate words corresponds to the intended speech in a given time interval, and (iii) synthesizing audible speech responsively to the extracted phonemes, such that in the at least some of the time intervals, the audible speech is synthesized by mixing the multiple candidate phonemes responsively to the respective probabilities.

[0044] As described below, to perform step (i) the processor may run an image processing software, and to perform step (ii) the processor may run a neural network. To perform step (iii) the processor may use a voice synthesizer.

[0045] The disclosed technique can be used, using the sensed movements of at least one of the target regions of skin on a face of a test subject, measuring the amount of neural activity, to indicate an intent of speech by the test subject even before such speech ever occurred. Finally, in another embodiment, the disclosed technique improves audio quality, for example, of conversations made by mobile telephones in loud public spaces, by cleaning (e.g., removing background signals) from audio.

System Description

[0046] FIG. 1 is a schematic pictorial illustration of a system **18** for silent speech sensing, in accordance with an embodiment of the invention. System **18** is based on a sensing device **20**, in which a bracket, in the form of an ear clip **22**, fits over the ear of a user **24** of the device. An earphone **26** attached to ear clip **22** fits into the user's ear. An optical sensing head **28** is connected by a short arm **30** to ear clip **22** (e.g., an AirPods) and thus is held in a location in proximity to the user's face. In the pictured embodiment, device **20** has the form and appearance of a clip-on headphone, with the optical sensing head in place of (or in addition to) the microphone.

[0047] Details of device **20**, such as of interface and processing circuitries comprised in device **20**, are described in the above-mentioned International Patent Application PCT/IB2022/054527.

[0048] Optical sensing head **28** directs one or more beams of coherent light toward different, respective locations on the face of user **24**, thus creating an array of spots **32** extending over an area **34** of the face (and specifically over the user's cheek). In the present embodiment, optical sensing head **28** does not contact the user's skin at all, but rather is held at a certain distance from the skin surface. Typically, this distance is at least 5 mm, and it may be even greater, for example at least 1 cm or even 2 cm or more from the skin surface. To enable sensing the motion of different parts of the facial muscles, the area **34** covered by spots **32** and sensed by optical sensing head **28** typically has an extent of at least 1 cm.^{sup.2}; and larger areas, for example at least 2 cm.^{sup.2} or even greater than 4 cm.^{sup.2}, can be advantageous.

[0049] Optical sensing head **28** senses the coherent light that is reflected from spots **32** the face and outputs a signal in response to the detected light. Specifically, optical sensing head **28** senses the secondary coherent light patterns that arise due to reflection of the coherent light from each of spots **32** within its field of view. To cover a sufficiently large area **34**, this field of view typically has a wide angular extent, typically with an angular width of at least 60°, or possibly 70° or even 90° or more. Within this field of view, device **20** may sense and process the signals due to the secondary coherent light patterns of all of spots **32** or of only a certain subset of spots **32**. For example, device **20** may select a subset of the spots that is found to give the largest amount of useful and reliable information with respect to the relevant movements of the skin surface of user **24**.

[0050] Within system **18**, processing circuitry processes the signal that is output by optical sensing

head **28** to generate a speech output. As noted earlier, the processing circuitry is capable of sensing movements of the skin of user **22** and generating the speech output, even without vocalization of the speech or utterance of any other sounds by user **22**. The speech output may take the form of a synthesized audio signal or a textual transcription, or both. In that regard, the silent speech detection can be readily implemented as nerve-to-text application, such as, for example, directly transcribing silent speech into an email draft. The synthesized audio signal may be played back via the speaker in earphone **26** (and is useful in giving user **22** feedback with respect to the speech output). Additionally or alternatively, the synthesized audio signal may be transmitted over a network, for example via a communication link with a mobile communication device, such as a smartphone **36**. Typically, the synthesis is done at different times than a voiced utterance would happen. This timing can be shorter or longer, and the processor can find the timing difference. Such timing difference may be utilized, as an example, when the synthesized voice is ready earlier than the voiced utterance would happen, to provide a translation of the synthesized voice into another language, with the translated utterance outputted on the time the voiced utterance would.

[0051] The functions of the processing circuitry in system **18** may be carried out entirely within device **20**, or they may alternatively be distributed between device **20** and an external processor, such as a processor in smartphone **36** running suitable application software. For example, the processing circuitry within device **20** may digitize and encode the signals output by optical sensing head **28** and transmit the encoded signals over the communication link to smartphone **36**. This communication link may be wired or wireless, for example using the Bluetooth™ wireless interface provided by the smartphone. The processor in smartphone **36** processes the encoded signal in order to generate the speech output. Smartphone **36** may also access a server **38** over a data network, such as the Internet, in order to upload data and download software updates, for example. Details of the design and operation of the processing circuitry are described hereinbelow with reference to FIG. **3**.

[0052] In the pictured embodiment, device **20** also comprises a user control **35**, for example in the form of a push-button or proximity sensor, which is connected to ear clip **22**. User control **35** senses gestures performed by user, such as pressing on user control **35** or otherwise bringing the user's finger or hand into proximity with the user control. In response to the appropriate user gesture, the processing circuitry changes the operational state of device **20**. For example, user **24** may switch device **20** from an idle mode to an active mode in this fashion, and thus signal that the device should begin sensing and generating a speech output. This sort of switching is useful in conserving battery power in device **20**. Alternatively or additionally, other means may be applied in controlling the operational state of device **20** and reducing unnecessary power consumption, for example as described below with reference to FIG. **5**. Moreover, a processor of device **20** can automatically switch from idle mode to high power consumption mode based on differing trigger types, such as a sensed input (e.g., eye blinks or mouth slightly open, or a pre-set sequence of motions like tongue movement). Also, the user may activate the device, using, for example, a touch button on the device, or from an application in a mobile phone.

[0053] In an optional embodiment, a microphone (not shown), may be included, to sense sound uttered by user **24**, enabling user **22** to use device **20** as a conventional headphone when desired. Additionally or alternatively, the microphone may be used in conjunction with the silent speech sensing capabilities of device **20**. For example, the microphone may be used in a calibration procedure, in which optical sensing head **28** senses movement of the skin while user **22** utters certain phonemes or words. The processing circuitry may then compare the signal output by optical sensing head **28** to the sounds sensed by a microphone (not shown) in order to calibrate the optical sensing head. This calibration may include prompting user **22** to shift the position of optical sensing head **28** in order to align the optical components in the desired position relative to the user's cheek.

[0054] FIG. **2** is a schematic pictorial illustration of a silent speech sensing device **60**, in accordance with another embodiment of the invention. In this embodiment, ear clip **22** is integrated with or otherwise attached to a spectacle frame **62**. Nasal electrodes **64** and temporal electrodes **66** are attached to frame **62** and contact the user's skin surface. Electrodes **64** and **66** receive body surface

electromyogram (sEMG) signals, which provide additional information regarding the activation of the user's facial muscles. The processing circuitry in device **60** uses the electrical activity sensed by electrodes **64** and **66** together with the output signal from optical sensing head **28** in generating the speech output from device **60**.

[0055] Additionally or alternatively, device **60** includes one or more additional optical sensing heads **68**, similar to optical sensing head **28**, for sensing skin movements in other areas of the user's face, such as eye movement. These additional optical sensing heads may be used together with or instead of optical sensing head **28**.

[0056] FIG. **3** is a block diagram that schematically illustrates functional components of system **18** for silent speech sensing, in accordance with an embodiment of the invention. The pictured system is built around the components shown in FIG. **1**, including sensing device **20**, smartphone **36**, and server **38**. Alternatively, the functions illustrated in FIG. **3** and described below may be implemented and distributed differently among the components of the system. For example, some or all of the processing capabilities attributed to smartphone **36** may be implemented in sensing device; or the sensing capabilities of device **20** may be implemented in smartphone **36**.

[0057] Sensing device **20** transmits the encoded signals via a communication interface of the device, such as a Bluetooth interface, to a corresponding communication interface **77** in smartphone **36**. In the present embodiment, the encoded output signals from sensing device **20** are received in a memory **78** of smartphone **36** and processed by a speech generation application **80** running on the processor in smartphone **36**. Speech generation application **80** converts the features in the output signal to a sequence of words, in the form of text and/or an audio output signal. Communication interface **77** passes the audio output signal back to speaker **26** of sensing device **20** for playback to the user. The text and/or audio output from speech generation application **80** is also input to other applications **84**, such as voice and/or text communication applications, as well as a recording application. The communication applications communicate over a cellular or Wi-Fi network, for example, via a data communication interface **86**.

[0058] The encoding operations of by device **20** and speech generation application **80** are controlled by a local training interface **82**. For example, interface **82** may indicate to a processor of device **20** which temporal and spectral features to extract from the signals output by receiver module **48** and may provide speech generation application **80** with coefficients of a neural network, which converts the features to words. In the present example, speech generation application **80** implements an inference network, which finds the sequence of words having the highest probability of corresponding to the encoded signal features received from sensing device **20**. Local training interface **82** receives the coefficients of the inference network from server **38**, which may also update the coefficients periodically.

[0059] To generate local training instructions by training interface **82**, server **38** uses a data repository **88** containing coherent light (e.g., speckle) images and corresponding ground truth spoken words from a collection of training data **90**. Repository **88** also receives training data collected from sensing devices **20** in the field. For example, the training data may comprise signals collected from sensing devices **20** while users articulate certain sounds and words (possibly including both silent and vocalized speech). This combination of general training data **90** with personal training data received from the user of each sensing device **20** enables server **38** to derive optimal inference network coefficients for each user.

[0060] Server **38** applies image analysis tools **94** to extract features from the coherent light images in repository **88**. These image features are input as training data to a neural network **96**, together with a corresponding dictionary **104** of words and a language model **100**, which defines both the phonetic structure and syntactical rules of the specific language used in the training data. Neural network **96** generates optimal coefficients for an inference network **102**, which converts an input sequence of feature sets, which have been extracted from a corresponding sequence of coherent light measurements, into corresponding phonemes and ultimately into an output sequence of words. Server **38** downloads the coefficients of inference network **102** to smartphone **36** for used in speech

generation application **80**.

Method for Speech Sensing

[0061] FIG. **4** is a flow chart that schematically illustrates a method for silent speech sensing, in accordance with an embodiment of the invention. This method is described, for the sake of convenience and clarity, with reference to the elements of system **18**, as shown in FIGS. **1** and **4** and described above. Alternatively, the principles of this method may be applied in other system configurations, for example using sensing device **60** (FIG. **2**) or a sensing device that is integrated in a mobile communication device.

[0062] As long as user **24** is not speaking, sensing device **20** operates in a low-power idle mode in order to conserve power of its battery, at an idling step **410**. This mode may use a low frame rate, for example twenty frames/sec. While device **20** operates at this low frame rate, it processes the images to detect a movement of the face that is indicative of speech, at a motion detection step **112**. When such movement is detected, a processor of device **20** instructs to increase the frame rate, for example to the range of 100-200 frames/sec, to enable detection of changes in the secondary coherent light (e.g., speckle) patterns that occur due to silent speech, at an active capture step **414**. Alternatively or additionally, the increase the frame rate may follow instructions received from smartphone **36**.

[0063] A processor of device **20** then extracts features of optical coherent light pattern motion, at a feature extraction step **420**. Additionally or alternatively, the processor may extract other temporal and/or spectral features of the coherent light in the selected subset of spots. Device **20** conveys these features to speech generation application **80** (running on smartphone **36**), which inputs vectors of the feature values to the inference network **102** that was downloaded from server **38**, at a feature input step **422**.

[0064] Based on the sequence of feature vectors that is input to the inference network over time, speech generation application **80** outputs a stream of words, which are concatenated together into sentences, at a speech output step **424**. As noted earlier, the speech output is used to synthesize an audio signal, for playback via speaker **26**. Other applications **84** running on smartphone **36** post-process the speech and/or audio signal to record the corresponding text and/or to transmit speech or text data over a network, at a post-processing step **426**.

Deciphering of Detected Silent Speech

[0065] As described above, the deciphering of silent speech (i.e., analyzing acquired signals to extract words corresponding to an intended speech) is performed by a chain of software tools, such as image processing software (e.g., tool **94**), and an artificial neural network (ANN), such as NN **96**. The image processing software converts acquired signals into preprocessed signals, and the ANN specifies intended speech words contained in the preprocessed signals. This section provides examples of deciphering methods and software tools that the disclosed technique may use. It covers training and inference phases by an ANN (FIGS. **5** and **6**, respectively), as well as the preprocessing phase (FIG. **7**).

[0066] FIG. **5** is a flow chart that schematically illustrates a method for training an ANN to perform silent speech deciphering, in accordance with an embodiment of the invention. This method can be used to train, for example, two different ANN types: a classification neural network that eventually outputs words, and a sequence-to-sequence neural network which outputs a sentence (word sequence). The process begins in data uploading step **502**, with uploading from a memory of server **38** pre-processed training data, such as outputted by image analyses tool **94**, that was gathered from multiple reference human subjects, e.g., during a development.

[0067] The silent speech data is collected from a wide variety of people (people of varying ages, genders, ethnicities, physical disabilities, etc.). The number of examples required for learning and generalization is task dependent. For word/utterance prediction (within a closed group) at least several thousands of examples were gathered. For the task of word/phoneme sequence prediction, the dataset size is measured in hours, and several thousands of hours were gathered for transcription.

[0068] In data augmentation step **504**, a processor augments the image processed training data to get more artificial data for the training process. In particular, the input here is an image processed

secondary coherent light pattern, with some of the image processing steps described below. Step 504 of data augmentation may include the sub-steps of (i) Time dropout, where amplitudes at random time points are replaced by zeros, (ii) Frequency dropout—the signal is transformed into the frequency domain. Random frequency chunks are filtered out. (iii) Clipping, where the maximum amplitude of the signal at random time points is clamped. This adds a saturation effect to the data, (iv) Noise addition, where gaussian noise is added to the signal, and speed change, where the signal is resampled to achieve a slightly lower or slightly faster signal.

[0069] At features extraction step 506, the augmented dataset goes through the feature extraction module. In this step the processor computes time domain silent speech features. For this purpose, for example, each signal is split into low and high frequency components, $x_{\text{sub.low}}$ and $x_{\text{sub.high}}$, and windowed to create time frames, using a frame length of 27 ms and shift of 10 ms. For each frame we compute five time-domain features and nine frequency domain features, a total of 14 features per signal. The time-domain features are as follows:

[00001]

$$\left[\frac{1}{n} \cdot \text{Math.} (x_{\text{low}}[i])^2, \frac{1}{n} \cdot \text{Math.} x_{\text{low}}[i], \frac{1}{n} \cdot \text{Math.} (x_{\text{high}}[i])^2, \frac{1}{n} \cdot \text{Math.} x_{\text{high}}[i] \cdot \text{Math.}, \text{ZCR}(x_{\text{high}}) \right]$$

where ZCR is the zero-crossing rate. In addition, we use magnitude values from a 16-point short Fourier transform, i.e., frequency domain features. All features are normalized to zero mean unit variance.

[0070] For ANN training step 508, the processor split the data into training, validation and test sets. The training set is the data used to train the model. Hyperparameter tuning is done using the validation set, and final evaluation is done using the test set.

[0071] The model architecture is task dependent. Two different examples describe training two networks for two conceptually different tasks. First is the signal transcription, i.e., translating silent speech to text by word/phoneme/letter generation. This task is addressed by using a sequence-to-sequence model. The second task is word/utterance prediction, i.e., categorizing utterances uttered by users into a single category within a closed group. It is addressed by using a classification model.

[0072] The disclosed sequence-to-sequence model is composed of an encoder, which transforms the input signal into high level representations (embeddings), and a decoder, which produces linguistic outputs (i.e., characters or words) from the encoded representations. The input entering the encoder is a sequence of feature vectors, as described in the “feature extraction” module. It enters the first layer of the encoder—a temporal convolution layer, which down samples the data to achieve a good performance. The model may use an order of hundred such convolution layers.

[0073] Outputs from the temporal convolution layer at each time step are passed to three layers of bidirectional recurrent neural networks (RNN). The processor employs long short-term memory (LSTM) as units in each RNN layer. Each RNN state is a concatenation of the state of the forward RNN with the state of the backward RNN. The decoder RNN is initialized with the final state of the encoder RNN (concatenation of the final state of the forward encoder RNN with the first state of the backward encoder RNN). At each time step, it gets as input the preceding word, encoded one-hot and embedded in a 150-dimensional space with a fully connected layer. Its output is projected through a matrix into the space of words or phonemes (depending on the training data).

[0074] The sequence-to-sequence model conditions the next step prediction on the previous prediction. During learning, a log probability is maximized:

$$[00002] \max_i \text{Math.} \log P(y_i | x, y_{<i};)$$

where $y_{<i}$ is the ground truth of the previous prediction. The classification neural network is composed of the encoder as in the sequence-to-sequence network and an additional fully connected classification layer on top of the encoder output. The output is projected into the space of closed words and the scores are translated into probabilities for each word in the dictionary.

[0075] The results of the above entire procedure are two types of trained ANNs, expressed in computed coefficients for an inference network 102. The coefficients are stored (step 510) in a memory of server 38.

[0076] In day-to-day use, training interface **82** receives up to date coefficients of inference network **102** from server **38**, where server **38** may also update the coefficients periodically, the coefficients of inference network **102** are stored in a memory of earpiece device **20** or in memory **78** of smartphone **36**. First ANN task is the signal transcription, i.e., translating silent speech to text by word/phoneme/letter generation. The second ANN task is word/utterance prediction, i.e., categorizing utterances uttered by users into a single category within closed group. These networks are plugged in the system to work as part of it at he below described in FIG. **6**.

[0077] Finally, the training session is used for optimizing a selection and priority of locations of secondary coherent light on the face of user **24** to analyze. In selection updating step **512**, the processor updates a list and order of use of such locations.

[0078] FIG. **6** is a flow chart that schematically illustrates a method of using a trained ANN in inference to perform silent speech deciphering, in accordance with an embodiment of the invention. Such trained ANN may be inference network **102**. The process begins with, for example, the processor in smartphone **36** running suitable application software that uploads inference network **102**, at trained ANN uploading step **602**.

[0079] At a silent speech preprocessing step **604**, a processor of sensing device **20** receives silent speech signals and preprocess this using, for example, an image processing software included in device **20**.

[0080] At silent speech feature extraction step **606**, the processor of sensing device **20** extracts from the preprocessed silent speech signals silent speech features, as described in FIG. **7**.

[0081] At silent speech features receiving step **608**, smartphone **36** receives encoded signals via a communication interface **77**. In the present embodiment, the encoded signals of step **606** from sensing device **20** are received in a memory **78** of smartphone **36**.

[0082] At a silent speech inference step **610**, the extracted features are processed by a speech generation application **80** running on the processor in smartphone **36**. Speech generation application **80** runs inference network **102** to converts the features in the output signal to a sequence (**612**) of words. These words may be subsequently outputted in a form of text and/or an audio output signal (e.g., using voice synthesizer).

[0083] FIG. **7** is a flow chart that schematically illustrates a method of preprocessing silent speech sensed data in preparation for speech deciphering, in accordance with an embodiment of the invention. The process begins by a processor of device **20** receiving a frame from a camera of device **20** that captured at secondary coherent light reflections from the cheek skin area at high frame rate (e.g., 500 fps), at a frame receiving step **702**.

[0084] For each frame the raw image is transferred to an image processing algorithm that extracts the skin motion at a set of pre-selected locations on the user's face. The number of locations to inspect is an input to the algorithm. The locations on the skin that are extracted for coherent light processing are taken from a predetermined list that a processor uploads (**704**) from memory. The list specifies anatomical locations, for example: check above mouth, chin, mid-jaw, check below mouth, high cheek and back of check. Furthermore, the list is dynamically updated with more points on the face that are extracted during the training phase (in step **512** of FIG. **2**). The entire set of locations is ordered in descending order such that any subset of the list (in order) minimizes the word error rate (WER) with respect to the chosen number of locations that are inspected.

[0085] At a coherent light spot selection step **706**, the processor selects the locations to analyze according to the list provided in step **704**.

[0086] At a cropping step **708**, the processor crops each of the coherent light spots that were extracted in the frame around the coherent light spots, and the algorithm process the spot. Typically, the process of coherent light spot processing involves reducing by two order of magnitude a size of full frame image pixels (of ~1.5MP) that is taken with the camera, with a very short exposure. Exposure is dynamically set and adapted to be able to capture only coherent light reflections and not skin segments. For day light and green laser this is found to be around 1/4000 seconds. As the image is mostly empty (e.g., of cheek skin-being black regions) and includes the laser point that forms a

secondary coherent light pattern. In the preprocessing phase, the laser point (e.g., speckle) region is identified, and the image is cropped, so that the algorithms run only on this region. For example, the processor reduces the full image (1.5MP) to a 18KP image, which immediately accelerates processing time for the remainder of the algorithm.

[0087] The image processing algorithm may be executed on a CPU, GPU, or hardware logic that is embedded within or adjacent to the camera sensor chip—to be able to eliminate the flow of high bandwidth data in the system. This may reduce the total power consumption of the device. Moreover, the preprocessed data corresponds to the physical vibrations and movements of the user's facial skin (in each location), thus is at much lower bandwidth compared to the raw images—few hundred of samples per second for each location.

[0088] Once the image processor has identified the region of interest within the coherent light spot, it improves the image contrast, by removing noise using a threshold s to determine black pixels and then computes (710) a characteristic metric of the coherent light. such as scalar speckle energy measure, e.g., an average intensity. Step 710 includes steps described in box 711, which includes, after identifying the coherent light pattern at step 706 and cropping it at step 708, further reducing the image to a predefined fraction (e.g., $\frac{1}{3}$) of the radius of the coherent light spot (which amounts to reducing the aforementioned 18KP image to only 2KP, upon which the metric is calculated. e.g., as average intensity of the 2KP pixels.

[0089] Analyzing changes in time in the measure (e.g., in average speckle intensity) by the processor is one example of detection of changes in the secondary coherent light patterns. Alternatively, other metrics may be used such as the detection of specific coherent light patterns. Finally, a sequence of values of this scalar energy metric is calculated frame-by-frame and aggregated (712), giving a 1D temporal signal.

[0090] The 1D temporal signals are the preprocessed signals that are stored for use in silent speech deciphering, as described above in FIGS. 5 and 6.

[0091] The accuracy of the word detection process described above is optimized using a combination of the following concepts:

1) Personalized Algorithm Parameters.

[0092] During normal speech of the user, the system simultaneously samples the user's voice and the facial movements. Automatic speech recognition (ASR) and Natural Language Processing (NLP) algorithms are applied on the actual voice, and the outcome of these algorithms is used for optimizing the parameters of the motion to language algorithms. These parameters include the weights of the various neural networks, as well as the spatial distribution of laser beams for optimal performance. For subjects with speech disorders who have intact nervous systems and muscle fibers, transfer learning techniques can be used in applying results acquired from subjects who are capable of speech.

2) Limiting the Word Set.

[0093] Limiting the output of the algorithms to a pre-defined word set significantly increases the accuracy of word detection in cases of ambiguity—where 2 different words result in similar movements on the skin.

[0094] The used word set can be personalized over time, adjusting the dictionary to the actual words used by the specific user, with their respective frequency and context.

3) Context Optimized Word Set.

[0095] Including the context of the conversation in the input of the words and sentences extraction algorithms increases the accuracy by eliminating out-of-context options. The context of the conversation is understood by applying Automatic speech recognition (ASR) and Natural Language Processing (NLP) algorithms on the other side's voice.

Voice Synthesis

[0096] The information that is extracted from the inner/silent speech can be used in various ways, e.g., 1) human-machine communication (e.g., personal assistant/“Alexa” type devices) 2) human—human communication (e.g., phone calls).

[0097] For human-human communication the system generates a synthetic voice of the user based on

the inner speech and transmits this synthetic voice to the other side's device. Alternatively, human-human communication can be made via 3rd party applications such as instant messaging apps, in which case, the inner speech is converted into text and transmitted to the other side's device.

[0098] In many cases, speech recognition algorithms result with some ambiguous result. For example:

[0099] The user says the word “Justice”, the classification algorithm predicts that with 50% certainty the articulated word was “justice”, 30% “practice” and 20% “lattice”. In other implementations of NLP algorithms, the algorithm selects the right word based on the context of the whole sentence that in many cases is revealed only after the ambiguous word has been said.

[0100] In the case of Human-Human real time communication implantation, waiting for the sentence to complete before synthesizing the text to speech will result with a significant delay that might not be acceptable.

[0101] To solve this issue, the speech synthesizer is configured to generate an ambiguous output for an ambiguous input. Examples of ambiguous input are confusing words/phonemes. For example, the system may not fully determine whether the user said “down” or “town”. The unit of work in this case is therefore the sound (called “phoneme”) and not an entire word. In response, the system generates and transmits speech that is a mixture of the two candidate phonemes. In the above example, if the system is not certain whether the user said “down” or “town” then the resulting sound will be a mixture of “t” and “d” at the appropriate times.

[0102] To further illustrate the concept, for the above example constituting of the words “justice”, “practice” and “lattice,” the synthesizer will not send a clear “justice” word, although it is the option with the highest probability, but instead, the algorithm will create a sound that is a combination of articulated words weighted as being probably 50% “justice”, 30% “practice” and 20% “lattice”—same probabilities as the input. This implementation will transfer the ambiguity to the other person's brain—to be resolved at a later time, after the sentence is complete.

[0103] The algorithm for generating ambiguous words comprises two steps: [0104] a) Time scaling to make all words the same length of time [0105] b) Weighted average of the sound waveform

[0106] FIG. 8 is a flow chart that schematically illustrates a method of generating an ambiguous speech output for an ambiguous silent speech input, in accordance with an embodiment of the invention. The process begins at words generation step **802**, with multiple candidate words are extracted by speech generation application **80** together with respective probabilities that each of the candidate words corresponds to the intended speech in a given time interval.

[0107] Next, at words synthesizing step **804**, a processor synthesizes the extracted words into audio signals (e.g., 1D sound waveforms). At a time-scaling step **806**, the processor, such as a one of a voice synthesizer (e.g., one example of application **84**), time scales similar words to make all words sound over a same time duration within a given time interval. Finally, at sound mixing step **808**, the audible speech is synthesized into an ambiguous audio output by mixing the multiple words responsively to the respective probabilities (e.g., by the processor performing weighted summation of the sound waveform amplitudes with the probabilities that serve as respective weights).

[0108] It will be appreciated that the embodiments described above are cited by way of example, and that the present invention is not limited to what has been particularly shown and described hereinabove. Rather, the scope of the present invention includes both combinations and subcombinations of the various features described hereinabove, as well as variations and modifications thereof which would occur to persons skilled in the art upon reading the foregoing description and which are not disclosed in the prior art.

Claims

1-37. (canceled)

38. A system for synthesizing speech from minute facial skin movements, the system comprising: a housing; at least one sensor connected to the housing in a manner such that when the housing is worn,

the at least one sensor is held at a distance from a skin surface; and at least one processor configured to: control at least one light source to project a plurality of light spots on a region of a face; receive from the at least one sensor, reflection signals indicative of light reflected from at least a subset of the plurality of light spots projected on the region of the face; analyze the reflection signals to determine minute facial skin movements indicative of silent speech; based on the determined minute facial skin movements indicative of the silent speech, extract words corresponding to the silent speech; and synthesize the words corresponding to the silent speech into audio signals.

39. The system of claim 38, wherein the at least one processor is further configured to extract the words corresponding to the silent speech using training-derived data stored in memory.

40. The system of claim 38, further comprising an earphone and wherein the at least one processor is further configured output the audio signals for presentation via the earphone.

41. The system of claim 38, further comprising a communication interface configured to transmit the audio signals to a mobile communications device via a communication link.

42. The system of claim 41, wherein the at least one processor is further configured to transmit to the mobile communications device a synthetization of the silent speech during a phone call.

43. The system of claim 38, wherein synthesizing the words includes translating the words to a language other than a language of the silent speech.

44. The system of claim 38, wherein the at least one processor is configured to analyze changes in average speckle intensity in the reflection signals to determine the minute facial skin movements indicative of the silent speech.

45. The system of claim 38, wherein synthesizing the words corresponding to the silent speech into the audio signals includes cleaning background signals.

46. The system of claim 38, wherein extracting the words corresponding to the silent speech includes determining multiple candidate phonemes from the minute facial skin movements based on respective probabilities that the minute facial skin movements are associated with the multiple candidate phonemes.

47. The system of claim 46, wherein the at least one processor is configured to mix the multiple candidate phonemes based on their respective probabilities to generate an ambiguous audio output.

48. The system of claim 47, wherein the generated ambiguous audio output is associated with a same length of time as a corresponding set of the minute facial skin movements.

49. The system of claim 38, wherein the at least one processor is configured to determine the silent speech in an absence of vocalization of the words.

50. The system of claim 38, wherein the at least one light source, the at least one sensor, and the at least one processor are part of a wireless headphone.

51. The system of claim 38, wherein the at least one light source is configured to direct a plurality of beams of light toward different locations on the region of a face thus creating an array of spots.

52. The system of claim 38, wherein the processor is further configured to extract temporal features associated with light patterns from the reflection signals and determine the minute facial skin movements associated with silent speech based on the temporal features.

53. The system of claim 38, wherein the processor is further configured to extract spectral features associated with light patterns from the reflection signals and determine the minute facial skin movements associated with silent speech based on the spectral features.

54. The system of claim 38, wherein extracting the words corresponding to the silent speech involves determining conversation context.

55. The system of claim 38, wherein the at least one processor is configured to analyze light patterns that include secondary light patterns due to reflection of the light reflected from the face to determine the minute facial skin movements indicative of the silent speech.

56. The system of claim 38, wherein the housing is associated with an earphone and an arm extending therefrom and wherein the at least one light source is located in the arm.

57. A method for synthesizing speech from minute facial skin movements, the method comprising: controlling at least one light source to project a plurality of light spots on a region of a face; receiving

from at least one sensor, reflection signals indicative of light reflected from at least a subset of the plurality of light spots projected on the region of the face; analyzing the reflection signals to determine minute facial skin movements indicative of silent speech; based on the determined minute facial skin movements indicative of the silent speech, extracting words corresponding to the silent speech; and synthesizing the words corresponding to the silent speech into audio signals.
