

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250266075

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Prather; Matthew A. et al.

REFRESH OPERATION IN MULTI-DIE MEMORY

Abstract

Methods, apparatuses, and systems for staggering refresh operations to memory arrays in different dies of a three-dimensional stacked (3DS) memory device are described. A 3DS memory device may include one die or layer of that controls or regulates commands, including refresh commands, to other dies or layers of the memory device. For example, one die of the 3DS memory may delay a refresh command when issuing the multiple concurrent memory refreshes would cause some problematic performance condition, such as high peak current, within the memory device.

Inventors: Prather; Matthew A. (Boise, ID), Kinsley; Thomas H. (Boise, ID)

Applicant: Lodestar Licensing Group LLC (Evanston, IL)

Family ID: 1000008578137

Appl. No.: 19/200309

Filed: May 06, 2025

Related U.S. Application Data

parent US continuation 18373120 20230926 parent-grant-document US 12315547 child US 19200309

parent US continuation 17379422 20210719 parent-grant-document US 11783882 child US 18373120

parent US continuation 16562940 20190906 parent-grant-document US 11069394 child US 17379422

Publication Classification

Int. Cl.: G11C11/406 (20060101); G11C7/10 (20060101); G11C11/407 (20060101)

U.S. Cl.:

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION(S) [0001] This application is a continuation of U.S. patent application Ser. No. 18/373,120, filed Sep. 26, 2023, which is a continuation of U.S. patent application Ser. No. 17/379,422, filed Jul. 19, 2021, which is a continuation of U.S. patent application Ser. No. 16/562,940, filed Sep. 6, 2019; each of which is incorporated herein by reference in its entirety.

TECHNICAL FIELD

[0002] The present disclosure generally relates to semiconductor memory devices, and more particularly relates to systems and methods for refreshing multi-die memory.

BACKGROUND

[0003] Memory devices are widely used to store information related to various electronic devices such as computers, wireless communication devices, cameras, digital displays, and the like. Memory devices are frequently provided as internal, semiconductor, integrated circuits and/or external removable devices in computers or other electronic devices. There are many different types of memory, including volatile and non-volatile memory. Volatile memory, including random-access memory (RAM), static random-access memory (SRAM), dynamic random-access memory (DRAM), and synchronous dynamic random-access memory (SDRAM), among others, can require a source of applied power to maintain its data. Non-volatile memory, by contrast, can retain its stored data even when not externally powered. Non-volatile memory is available in a wide variety of technologies, including flash memory (e.g., NAND and NOR), phase change memory (PCM), ferroelectric random-access memory (FeRAM), resistive random-access memory (RRAM), and magnetic random-access memory (MRAM), among others. Improving memory devices, generally, can include increasing memory cell density, increasing read/write speeds or otherwise reducing operational latency, increasing reliability, increasing data retention, reducing power consumption, or reducing manufacturing costs, among other metrics.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0004] FIG. 1 is a representative block diagram schematically illustrating a memory device.

[0005] FIG. 2 is a representative block diagram schematically illustrating a three-dimensional stacked (3DS) memory device.

[0006] FIG. 3A is a representative plot illustrating example timing for non-overlapping refresh operations to DRAMs in different dies in a multi-die memory device.

[0007] FIG. 3B is a representative plot illustrating example current profiles corresponding to non-overlapping refresh operations to DRAMs in different dies in a multi-die memory device.

[0008] FIG. 4A is a representative plot illustrating example timing for overlapping refresh operations to DRAMs in different dies in a multi-die memory device.

[0009] FIG. 4B is a representative plot illustrating example current profiles corresponding to overlapping refresh operations to DRAMs in different dies in a multi-die memory device.

[0010] FIG. 5 is a representative flow illustrating an example method to retime refresh operations for memory arrays in DRAMs in different dies in a multi-die memory device.

[0011] FIG. 6 is a representative flow illustrating an example method to retime refresh operations for DRAMs in different dies in a multi-die memory device according to another embodiment.

[0012] FIG. 7A is a representative plot illustrating retimed refresh operations for DRAMs in different dies in a multi-die memory device.

[0013] FIG. 7B is a representative plot illustrating current profiles corresponding to retimed refresh operations for DRAMs in different dies in a multi-die memory device.

DETAILED DESCRIPTION

[0014] DRAM memories need to periodically refresh the memory arrays to preserve the information in the memory by reading the information from the arrays and immediately rewriting the read information to the same areas without modification. Because refresh is power intensive, having multiple refresh operations occurring simultaneously in a memory device can require more power than is normally budgeted for operation. A monolithic device (e.g., a memory device comprising a single silicon die) can be configured to handle this by staggering the activate commands for the rows being refreshed to keep the power consumption within budget. For a multi-chip device, however, this staggering approach can be problematic—a delayed activate command from one die may directly coincide in time with an activate command from another die, causing the power requirements to exceed the normal power budget, which can, for example, detrimentally increase the voltage noise experienced by the memory device (interfering with reliable data communication and/or storage), or exceed a power level the host system can provide (potentially causing catastrophic failure).

[0015] To address the foregoing problem, systems and methods for refreshing multi-die memories are disclosed herein. Concurrent refresh or simultaneous refresh operations to multiple dies in a three-dimensional stacked (3DS) memory devices are staggered by delaying subsequent refresh commands to a memory array of the 3DS memory device when a preceding refresh command to another memory array of the 3DS memory device is ongoing. According to one embodiment, the master device delays refresh commands targeted at a memory array in a die to ensure that there are no overlapping refresh operations that can lead to high instantaneous peak current demands from the 3DS memory device.

[0016] As will be discussed further below, elements of the presently disclosed technology described in the context of paragraphs and/or Figures can be combined with elements described in the context of other paragraphs and/or Figures. Furthermore, not all elements of the technology disclosed herein are required to practice the technology. Additionally, several details describing structures and/or processes that are well-known and often associated with memory devices, but that may unnecessarily obscure some significant aspects of the present technology, are not set forth in the following description for purposes of clarity. Moreover, although the following disclosure sets forth several embodiments of the present technology, several other embodiments of the technology have different configurations or different components than those described in this section. As such, the present technology can have other embodiments, with additional elements and/or without several of the elements described below with reference to the figures.

[0017] FIG. 1 is a block diagram schematically illustrating a memory device **100** in accordance with an embodiment of the present technology. The memory device **100** can include an array of memory cells, such as memory array **150**. The memory array **150** can include a plurality of banks, and each bank can include a plurality of word lines (WL), a plurality of bit lines (BL), and a plurality of memory cells arranged at intersections of the word lines and the bit lines. Memory cells can include any one of different memory media types, including capacitive, magnetoresistive, ferroelectric, phase change, or the like. The selection of a word line WL can be performed by a row decoder **140**, and the selection of a bit line BL can be performed by a column decoder **145**. Sense amplifiers (SAMP) can be provided for corresponding bit lines BL and connected to at least one respective local I/O line pair (LIOT/B), which can in turn be coupled to at least respective one main I/O line pair (MIOT/B), via transfer gates (TG), which can function as switches. The memory array **150** can also include plate lines and corresponding circuitry for managing their operation.

[0018] The memory device **100** can employ a plurality of external terminals that include command

and address terminal (C/A) coupled to a command bus and an address bus to receive command signals (CMD) and address signals (ADDR). The memory device can further include a chip select terminal to receive a chip select signal CS, a chip identifier terminal to receive a chip ID signal CHIP_ID, clock terminals to receive clock signals CK and CKF, data terminals DQ and DQS, power supply terminals VDD, VSS, VDDQ, and VPP.

[0019] The command and address terminal (C/A) can be supplied with an address signal and a bank address signal from outside. The address signal and the bank address signal supplied to the address terminals can be transferred, via a command/address input circuit **105**, to an address decoder **110**. The address decoder **110** can receive the address signals and supply a decoded row address signal (XADD) to the row decoder **140**, and a decoded column address signal (YADD) to the column decoder **145**. The address decoder **110** can also receive the bank address signal (BADD) and supply the bank address signal to both the row decoder **140** and the column decoder **145**.

[0020] The command/address input circuit **105** can be supplied with command and address signals C/A, chip identifier signals CHIP_ID, and chip selection signals CS, from a memory controller, host, CPU, or SOC. In some embodiments, CHIP_ID is part of the C/A bus and in other embodiments CHIP_ID is a standalone input. The command signals can represent various memory commands from the memory controller (e.g., including access commands, which can include read commands and write commands and commands to refresh the memory array). The select signal CS and chip identifier CHIP_ID signals can be used to select the memory device **100** to respond to commands and addresses provided to the command and address terminal. When an active CS and CHIP_ID signal is provided to the memory device **100**, the commands and addresses can be decoded and memory operations can be performed. The command signals CMD can be provided as internal command signals ICMD to a command decoder **115** via the command/address input circuit **105**. The command decoder **115** can include circuits to decode the internal command signals ICMD to generate various internal signals and commands for performing memory operations. The command decoder **115** can further include one or more registers **117** for tracking various counts or values. The memory device **100** can also include a refresh controller **112** to control the refresh of the memory array **150** based on commands from the command decoder **115**. The refresh controller **112** can also include registers (not shown in FIG. **1**) for tracking various counts or values (e.g., counts of refresh commands received by the memory device **100** or self-refresh operations performed by the memory device **100**). It will be appreciated that the various partitions and functional blocks described herein are merely illustrative to enable an understanding of the disclosed technology and memory device **100** can be implemented with functionalities residing in different blocks and different logical partitions.

[0021] The power supply terminals may be supplied with power supply potentials VDD and VSS. These power supply potentials VDD and VSS can be supplied to an internal voltage regulator/generator circuit **170**. The internal voltage regulator circuit **170** can generate various internal potentials VPP, VOD, VARY, VPERI, and the like based on the power supply potentials VDD and VSS. The internal potential VPP can be used in the row decoder **140**, the internal potentials VOD and VARY can be used in the sense amplifiers included in the memory array **150**, and the internal potential VPERI can be used in many other circuit blocks.

[0022] The power supply terminal may also be supplied with power supply potential VDDQ. The power supply potential VDDQ can be supplied to the input/output circuit **160** together with the power supply potential VSS. The power supply potential VDDQ can be the same potential as the power supply potential VDD in an embodiment of the present technology. The power supply potential VDDQ can be a different potential from the power supply potential VDD in another embodiment of the present technology. However, the dedicated power supply potential VDDQ can be used for the input/output circuit **160** so that power supply noise generated by the input/output circuit **160** does not propagate to the other circuit blocks.

[0023] Input buffers included in the clock input circuit **120** can receive the external clock signals.

For example, an input buffer can receive the CK and CKF signals. The clock input circuit **120** can receive the external clock signals to generate internal clock signals ICLK. The internal clock signals ICLK can be supplied to an internal clock circuit **130**. The internal clock circuit **130** can provide various phase and frequency controlled internal clock signal based on the received internal clock signals ICLK. For example, the internal clock circuit **130** can include a clock path (not shown in FIG. 1) that receives the internal clock signal ICLK and provides various clock signals to the command decoder **115**. The internal clock circuit **130** can further provide input/output (IO) clock signals. The IO clock signals can be supplied to the input/output circuit **160** and can be used as a timing signal for determining an output timing of read data and the input timing of write data. The IO clock signals can be provided at multiple clock frequencies so that data can be output from and input to the memory device **100** at different data rates. A higher clock frequency may be desirable when high memory speed is desired. A lower clock frequency may be desirable when lower power consumption is desired. The internal clock signals ICLK can also be supplied to a timing generator **135** and thus various internal clock signals can be generated.

[0024] To refresh all the DRAM cells or memory device **100** within the allotted refresh period, a Refresh command (REF) must be issued at the average periodic refresh interval time (tREFI). For example, in some embodiments when CS_n (inverse of CS in FIG. 1), RAS_n/A16 and CAS_n/A15 are held Low, and WE_n/A14 and ACT_n (not shown in FIG. 1) is held High at the rising edge of the clock (CK), memory device **100** can enter a Refresh cycle. In some embodiments, the Refresh cycle is entered by decoding the C/A bus without use of RAS, CAS, and WE inputs. All banks of the memory device **100** must be precharged and idle for a minimum of the precharge time tRP (min) before the Refresh Command (REF) can be applied. The refresh addressing is generated by the internal refresh controller **112**. This makes the address bits inconsequential (e.g., they may be assigned a logic value of “Don't Care”) during a Refresh command. An internal address counter supplies the addresses during the refresh cycle. No control of the external address bus is required once this cycle has started. When the refresh cycle has completed, all banks of the memory device **100** will be in the precharged (idle) state. A delay between the Refresh Command and the next valid command, except DES, must be greater than or equal to the minimum Refresh cycle time tRFC (min) where the tRFC timing parameter depends on memory density.

[0025] In general, a Refresh command (REF) needs to be issued to the memory device **100** regularly every tREFI interval. To allow for improved efficiency in scheduling and switching between tasks, some flexibility in the absolute refresh interval is provided for postponing and pulling-in refresh command. For example, a maximum of 8 Refresh commands can be postponed when memory device **100** is in 1× refresh mode; for 2×/4× refresh mode, 16/32 Refresh commands can be postponed respectively during operation of the memory device **100**, meaning that at no point in time more than a total of 8, 16, 32 Refresh commands are allowed to be postponed for 1×, 2×, 4× Refresh mode respectively. When 8 Refresh commands are postponed in a row, the resulting maximum interval between the surrounding Refresh commands is limited to 9×tREFI. Additionally, a maximum of 8 additional Refresh commands can be issued in advance (“pulled in”) in 1× refresh mode and for 2×/4× refresh mode, 16/32 Refresh commands can be pulled in respectively, with each one reducing the number of regular Refresh commands required later by one. Pulling in more than Aug. 16, 1932 Refresh commands in advance does not further reduce the number of regular Refresh commands required later. The resulting maximum interval between two surrounding Refresh commands is limited to 9× tREFI where 8 refresh commands are pulled in.

[0026] The memory device **100** can be connected to any one of several electronic devices capable of utilizing memory for the temporary or persistent storage of information, or a component thereof. For example, a host device of memory device **100** can be a computing device such as a desktop or portable computer, a server, a hand-held device (e.g., a mobile phone, a tablet, a digital reader, a digital media player), or some component thereof (e.g., a central processing unit, a co-processor, a

dedicated memory controller, etc.). The host device can be a networking device (e.g., a switch, a router, etc.) or a recorder of digital images, audio and/or video, a vehicle, an appliance, a toy, or any one of several other products. In one embodiment, the host device can be connected directly to memory device **100**, although in other embodiments, the host device can be indirectly connected to memory device (e.g., over a networked connection or through intermediary devices).

[0027] FIG. **2** is a representative block diagram **200** schematically illustrating a three-dimensional stacked (3DS) memory device **200**. The 3DS memory device **200** comprises several stacked memory devices, e.g., memory devices **100a**, **100b**, and **100c** interconnected through, for example, through-silicon via (TSV) or through-chip via interconnect. In the 3DS memory device **200**, each of the component memory devices **100a**, **100b**, and **100c** can be configured to operate as a master memory device (e.g., memory device **100a**) or as slave memory devices (e.g., memory devices **100b** and **100c**). Only the master memory device **100a** interfaces with the external memory controller (e.g., memory controller **210**) or host controller; the slave devices are controlled by the memory controller **210** through the master device **100a**. That is, memory controller instructions to the slave memory devices **100b** and **100c** pass through and are “filtered” by the master memory device **100a** (e.g., memory devices of a slave type receive commands or control information from a memory devices of a master type). In some embodiments, both the master device (e.g., master device **100a**) and the slave devices (e.g., slave devices **100b** and **100c**) have similar hardware components except that the slave devices are configured to disable “master” functionality, e.g., via fuse configuration, packing bonding, or by programming via for example mode registers. In other embodiments, only master devices include hardware for “master functionality,” e.g., hardware to interface memory controller **210**, but the slave devices exclude the master logic and circuitry altogether.

[0028] Each of the master and slave devices in 3DS memory device **200** includes a memory array, for example, memory array **150a** in the master device **100a**, and memory arrays **150b** and **150c** in memory devices **100b**, and **100c**, respectively. As discussed above in relation to FIG. **1**, DRAM memory arrays such as memory arrays **150a**, **150b**, and **150c** need to be periodically refreshed to preserve the information in the memory arrays. To refresh the memory arrays **150a**, **150b**, and **150c**, a Refresh command (REF) must be issued at an average periodic refresh interval time (tREFI) to each memory device **100a**, **100b**, and **100c**. Memory controller **210** issues the REF command for each of the memory devices **100a**, **100b**, and **100c**, but the REF commands for slave devices **100b** and **100c** are further retimed by the master device **100a** as described below.

[0029] The master device **100a** receives a REF command from the memory controller **210**. In some embodiments, the REF cycle is entered when CS_n (inverse of CS in FIG. **1**), RAS_n/A₁₆ and CAS_n/A₁₅ are held Low, and WE_n/A₁₄ and ACT_n (not shown in FIG. **1**) is held High at the rising edge of the clock (CK). In other embodiments, the Refresh cycle can be entered by decoding the C/A bus without use of RAS, CAS, and WE inputs. The master refresh controller **112a** generates the addresses and timing to refresh memory array **150a**. Additionally, the master device **100a** receives REF commands targeted to slave devices **100b** and **100c** (e.g., where CS_n chipselect selects slave device **100b** or **100c**). As described further below, master device **100a** determines whether to immediately send the REF command to the targeted slave device or whether and to what extent to delay the REF commands to the targeted slave device. In some embodiments, the master device **100a** delays the REF command for a certain period before refreshing the targeted memory array (the targeted memory array can be the master array **150a** or any of the slave arrays **150b** and **150c**). For example, if the refresh is targeted to a slave device, after the delay elapses, the master device **100a** sends the REF command to the individual slave device where each slave devices refresh controller (e.g., refresh controller **112b** corresponding to slave device **100b**, and refresh controller **112c** corresponding to slave device **100c**, or other control logic in slave devices **100b** and **100c** configured to receive refresh commands from the master memory device and to initiate and/or control refresh operations to slave memory arrays) controls the timing and

addressing of the refresh command. In other embodiments, the master refresh controller **112a** generates the refresh addressing and after the delay elapses provides the addressing to slave device refresh controllers, e.g., via path **225**. Additionally, in other embodiments, the master device **100a** includes the master/slave logic for the 3DS stack including the refresh controller for the master and slave devices (i.e., slave refresh controller **112b** and **112c** are contained in the master device **100a**). As described above, when the master device **100a** receives REF commands targeted to slave devices, the master device **100a** determines whether to immediately start a refresh operation to the targeted slave device or whether and to what extent to delay the start of the refresh operation to the targeted slave device.

[0030] FIG. **3A** is a representative plot illustrating example timing for non-overlapping refresh operations to memory arrays in a monolithic device (e.g., memory device **100**). A refresh operation involves activating and precharging tens to hundreds of rows in a short amount of time. Each activate and precharge is a comparatively high current operation which presents challenges in power distribution and noise immunity. On a monolithic device such as memory device **100**, the refresh controller (e.g., refresh controller **112** in FIG. **1**) manages the current spikes by staggering the internal timing of the activates to reduce the peak current demanded by the circuit. For example, memory controller **112** can refresh certain rows using timing **310a**, and refresh other rows using timing **310b**. The non-overlapping refresh operations based on timing **310a** and **310b** is achieved by staggering the internal timing of the activate commands to reduce the peak current demanded by the memory circuit. That is, at time **315a** and **317a** when the refresh to the first set of rows begins and ends (REF1), the second set of rows are not being refreshed; and at time **315b** and **317b** when the refresh to the second set of rows is ongoing (REF2), the first set of rows are not being refreshed. This non-overlapping timing in refresh operations results in non-overlapping peak currents as illustrated in FIG. **3B**, a representative plot of the current profile, e.g., average current drawn during a refresh (IDD5), for periodic and repetitive refreshes. That is, at time **316a** between times **315a** and **317a**, a peak current **350a** results from the refresh operation on the first set of memory rows; at a time **316b**, different from time **316a**, a peak current **350b** results from the refresh operation on the second set of memory rows. Because times **316a** and **316b** are sufficiently non-overlapping, peak currents **350a** and **350b** are not additive as would be the case if the refresh timing to the first and second set of rows overlapped.

[0031] FIG. **4A** is a representative plot illustrating example timing for overlapping refresh operations to memory arrays in different dies in a multi-die memory device (e.g., memory device **200**). For example, refresh timing **415a** corresponds to refresh of memory array **150a** in master memory device **100a**; refresh timing **417a** corresponds to refresh of memory array **150b** in slave memory device **100b**; and, refresh timing **419a** corresponds to refresh of memory array **150c** in slave memory device **100c**. At time **420a**, refresh to rows in memory arrays **150a**, **150b**, and **150c** is concurrently active which leads to a large peak current **450** around time **420b** as illustrated in FIG. **4B** (time **420b** corresponding to the high peak current falls around time **420a** when memory arrays in different dies of the 3DS memory are simultaneously refreshing). For example, if each 3DS die/rank (e.g., memory devices **100a**, **100b**, and **100c** in the 3DS device **200** of FIG. **2**) must complete its refresh operation in 295 ns but can be issued a REF command every 95 ns, there can be cases when 3 dies or logical ranks are in refresh simultaneously as in period **420a**.

[0032] FIG. **4B** illustrates the individual current profiles (e.g., IDD5 average refresh currents) for each refresh operation described in relation to FIG. **4A**; the current profile **415b** from refreshing rows in memory array **150a** based on refresh timing **415a**; the current profile **417b** from refreshing rows in memory array **150b** based on refresh timing **417a**; and the current profile **419b** from refreshing rows in memory array **150c** based on refresh timing **419a**. As seen in FIG. **4B**, the aggregate instantaneous peak current **450** for the 3DS memory can be higher than the peak currents **415b**, **417b**, and **419b** to the separate memory arrays in the individual dies when the individual refresh cycles overlap as in period **420a**. The high aggregate peak current **450** can affect data

retention margins, timing circuits, signal integrity, and can require more power than the system can provide.

[0033] FIG. 5 is a representative flow 500 illustrating an example method to retime refresh operations for DRAMs in different dies in a multi-die memory device (e.g., memory device 200). At block 510, the master DRAM (e.g., master device 100a in FIG. 2) receives a refresh command from the host, CPU or memory controller (e.g., memory controller 210 in FIG. 2) directed to a first memory array (e.g., any one of the memory arrays 150a, 150b, or 150c in FIG. 2). At block 520, the master device 100a receives a second refresh command to a different memory array in the 3DS stack. At block 530, the master device 100a determines if timely executing the second refresh command would overlap with execution of the first refresh command. That is, at block 530 the master device determines if executing the second refresh command within a time required by the specifications for the maximum refresh interval (e.g., tREFI) for the memory array relative to when the row of the memory array was last refreshed will result in the refresh overlapping with an ongoing refresh to a memory array of another die. For example, memory device 100a, on receiving the first refresh command to the first memory array at block 510 can commence or start execution of the refresh command to the targeted memory array and concurrently start a timer or counter. On receiving the second refresh command at block 520, memory device 100a can determine if the elapsed time or count since the start of the first refresh command is such that immediately allowing execution of the second refresh command would lead to overlap in refresh commands to different dies in the 3DS stack (i.e., rows in different arrays in different stacks would be simultaneously refreshing and hence consuming large amounts of peak current). In some embodiments, a free running timer or counter can be used to determine the relative times when refresh commands are received from the memory controller and issued by the master device and can be used to count the number of clock cycle delays required to delay execution of the received refresh commands.

[0034] At block 540, if allowing the execution of the second refresh command would lead to overlapping refresh operations, memory device 100a delays the execution of the second refresh command to the second memory array. In some embodiments, the amount of delay is based on when the second refresh command is received relative to when the execution of the first refresh command started and how long the refresh takes and is set such that execution of the second refresh command does not overlap with the execution of the first refresh command while ensuring that the delay is not too large to violate the maximum memory refresh period required to preserve the data in the memory array. For example, the master device 100a can determine the maximum amount of delay for the next refresh operation that is possible based on when the last previous refresh operation of the same rows of the same memory array occurred and further based on the required periodicity of row refresh required. The memory device 100a can also determine the minimum delay required to ensure no overlap between the execution of the second refresh and any ongoing refreshes. In some embodiments, the minimum delay is based at least in part on an operating characteristic of the 3DS memory device; or on a user-defined value (e.g., a value selected by a user in a programmable register); or on a multiple of a time period associated with refresh operations (e.g., if refresh operations take a time n for a given memory device, the minimum delay can be based on a multiple of n); or on the minimum delay that results in peak currents below a predetermined current threshold (e.g., where different preconfigurable amounts of delays result in different peak currents and a delay can be selected based on a minimum peak current target); or an industry specification (e.g., an industry specification based on one or more of the criteria above); or any combination thereof. Memory device 100a can select a delay value between the computed maximum and minimum delay (i.e., can select the maximum delay, the minimum delay, or a delay greater than the minimum and less than the maximum delay). In some embodiments, the delay is preconfigured or predetermined based on, for example, programmable register values, or based on some predetermined relationship based on other parameters and register values, or based on a hardcoded delay value (e.g., based on a fuse configuration, a logic circuit configuration, or other

static/fixed signal determined at the time of manufacture or power-on). Additional details on how memory device **100a** determines the amount of delay that is suitable is provided below in relation to FIG. 6.

[0035] FIG. 6 is a representative flow **600** illustrating an example method to retime refresh operations for DRAMs in different dies in a multi-die memory device (e.g., memory device **200**) according to another embodiment. At block **610**, the master DRAM device (e.g., master device **100a** in FIG. 2), detects a first refresh command to a first memory array. For example, the first refresh command can be a command to refresh rows of memory array **150a**, or **150b**, or **150c** in FIG. 2. At block **620**, the master device **100a** issues the refresh command to the targeted memory array. For example, if the refresh command is targeted to rows in memory array **150a** of master device **100a**, master refresh controller **112a** can generate the row addresses to refresh and can control the refresh of the memory array **150a**. If, on the other hand, the refresh command is targeted to rows in memory arrays **150b** or **150c** of slave devices **100b** or **100c**, the master device can instruct the slave refresh controllers **112b** or **112c** to start refresh of the targeted rows of memory arrays **150b** or **150c**.

[0036] At block **630**, the master device **100a** determines a time, t_1 , when the first refresh command was issued (or when the first refresh operation commenced or was issued). Master device **100a** can start a counter, zero out a running counter, or otherwise sample a time value coincident with issuing the first refresh command at block **620**. Establishing the time that the first refresh command is issued, enables master device **100a** determine the amount of time that has elapsed between the issuing of that refresh command and subsequent refresh commands. Knowing the elapsed time enables the master device **100a** to determine if subsequent refresh commands would interfere with ongoing refresh operations. The master device **100a** makes this determination on detecting a subsequent refresh command to a different memory array in block **640** (i.e., a memory array in another 3DS die in the stack), determining at block **650** a time, t_2 , when this subsequent refresh command is received from the memory controller or host, and determining, at block **660**, if the elapsed time is less than a predetermined time threshold, $t_{\text{threshold}}$. For example, if at block **630**, the time for issuing the first refresh command is determined by starting a counter with count values calibrated to denote time in ns, master device **100a** can set t_1 to 0 ns. Then at block **650**, master device **100a** can determine the current counter value, and determine a time of t_2 ns since the counter was started. At block **660**, master device **100a** can then determine if t_2 ns is less than a threshold delay value. The threshold delay value, $t_{\text{threshold}}$, is set to ensure no overlap between the first refresh and subsequent refreshes while ensuring that minimum refresh interval (t_{REFI}) requirement is not violated. For example, if $t_{\text{threshold}}$ is configured to 100 ns and t_2 is determined in block **650** to be 50 ns, the master device **100a** can determine at block **660** that $t_2 - t_1 < 100$ ns and delay the issuance of the second refresh command as described further below.

[0037] If at block **660**, master device **100a** determines that the subsequent refresh command has come in too soon (i.e., that $t_2 - t_1 < t_{\text{threshold}}$), master device **100a** delays, at block **670**, the issuance of the subsequent refresh command at block **680**. On the other hand, if at block **660** master device **100a** determines that there is sufficient delay between when the first refresh command was issued and when the second refresh command is received (i.e., that $t_2 - t_1 > t_{\text{threshold}}$), it issues, at block **680**, the second refresh command to the targeted memory array without further delaying the refresh command.

[0038] At block **670**, the amount of delay can be a predetermined, calculated, or configurable number of clock cycles or time intervals and determines when master device **100a** issues the received refresh command either targeted to the master device **100a**'s memory array **150a** or to slave devices **100b** or **100c**'s memory arrays **150b** or **150c**, respectively. In some embodiments, the delay can be made up of a fixed/configurable delay t_{dly} plus a delay corresponding to how much sooner the subsequent refresh command was received relative to the predetermined or configurable threshold $t_{\text{threshold}}$. For example, if t_{dly} is set to 2 ns (i.e., the refresh command to the second

memory array must lag the completion of the execution of the first refresh command to the first memory array by 2 ns or the preconfigured time after completion of the first refresh operation to start the second refresh operation is 2 ns), and $t_{\text{threshold}}$ is 100 ns, and the second refresh command is determined (at block **650**) to have been received 50 ns after the first refresh command was issued at block **620** (i.e., $t_2=50$ ns, $t_1=0$ ns), master device **100a** can delay issuing the second refresh command to the targeted memory by 52 ns ($t_{\text{dly}} + (t_{\text{threshold}} - (t_2 - t_1))$). It will be appreciated that the values in the example above are merely illustrative and different values for t_{dly} and $t_{\text{threshold}}$ can be selected that achieve the goal of ensuring that the subsequent refresh operation will not overlap in time with the first refresh operation (which leads to simultaneous high current demand from shared power rails). In some embodiments, the delay can be based at least in part on an operating characteristic of the 3DS memory device (e.g., peak currents), or can be a user defined value, or can be an industry specification, or any combination thereof. At block **670**, master device **100a** can execute the delay by, for example, starting a timer or counter (e.g., counting a number of clocks corresponding to the determined delay).

[0039] In some embodiments, the master device **100a** can also delay internal activate or precharge for master or slave DRAMs to delay the start of refresh to those DRAMS. In other embodiments, the amount of delay incorporates the time required to refresh the memory array such that a subsequent refresh to another memory array is not started until the preceding refresh is complete. For example, for a preconfigured minimum non-overlap of 2 ns, memory device **100a** can wait for the completion of the first refresh operation, wait 2 ns, and start the execution of the second refresh operation. In other embodiments, the amount of delay is based on target peak currents from overlapped refresh operations. That is, the delay is chosen to keep the peak aggregate current below a predetermined or preconfigured current threshold. In these embodiments, some amount of overlapping refresh operations can be allowed provided that the resulting peak refresh currents for the 3DS memory does not exceed some threshold peak current. At block **680**, when the master device **100a** issues or allows the execution of the refresh command, it can stop and reset the delay counter and start a new counter or timer corresponding to when it issued the second refresh command. In this case, the functions in blocks **640-680** are repeated when subsequent refresh commands (e.g., third, fourth, etc. refresh commands) are received by master device **100a**.

[0040] FIG. 7A is a representative plot **700a** illustrating retimed refresh operations for DRAMs in different dies in a multi-die memory device (e.g., memory device **200**). Curves **720a**, **722a**, and **724a** represents the retimed refresh timing to a first, second, and third memory array of different memory devices of the 3DS stack. For example, curve **720a** can represent the refresh timing of master device **100a**'s memory array **150a**, curve **722a** can represent the refresh timing of slave device **100b**'s memory array **150b**, and curve **724a** can represent the refresh timing of slave device **100c**'s memory array **150c**. Notably, plot **700a** illustrates the example where the refresh to memory array **150b** in accordance with timing **722a** is delayed by a time **710** from the completion of refresh to memory array **150a**; refresh to memory array **150c** according to timing **724a** is delayed by time **71** from the completion of refresh to memory array **150b**, etc. The amount of delays **710** and **711** can be determined as described above in relation to block **670** of FIG. 6. That is, if curve **720** represents the second/subsequent refresh command detected in block **640** (FIG. 6), because immediately issuing the refresh command would have resulted in an overlap of refresh operations (and hence higher peak currents), at block **670** (FIG. 6) the timing **720** is delayed by a time delay **709** to timing **722a**. For example, if the threshold time, $t_{\text{threshold}}=t_3$ (in FIG. 7A), and time **710** ($t_4 - t_3$) is the preconfigured delay, t_{dly} , described in relation to FIG. 6, the delay **709** is equal to $t_{\text{dly}} + (t_{\text{threshold}} - (t_2 - t_1))$ or $(t_4 - t_3) + (t_3 - (t_2 - t_1))$. In some embodiments, the amount of delay is fixed and selected large enough to always allow for no overlap (or overlap below a predetermined minimum overlap time) between subsequent memory refresh commands. For example, if a memory refresh command typically takes x ns to complete, a fixed delay of y ns, where $y > x$ can ensure that there is never an overlap of refresh commands to different memory arrays in the 3DS memory

device. In other embodiments, the amount of delay is variable based on the time when the first refresh operation was started, a time when the first refresh operation is expected to complete, and further base on a preconfigured or predetermined static time indicating when, after completion of the first refresh operation, the second refresh operation can begin.

[0041] FIG. 7B is a representative plot illustrating current profiles corresponding to retimed refresh operations of FIG. 7A. Current profile **720b** corresponds to refresh timing **720a** (for refresh operations to a first memory device in the 3DS stack); current profile **722b** corresponds to refresh timing **722a** (for refresh operations to a second memory device in the 3DS stack); and current profile **724b** corresponds to refresh timing **724a** (for refresh operations to the first memory device in the 3DS stack). Current profile **730** represents the aggregate current consumed by the 3DS stack device (e.g., memory device **200** of FIG. 2). For example, this can be the current drawn from shared power rails. As seen in FIG. 7B, because there is a delay between refresh operations of different memory arrays in the stack, the peak currents are non-overlapping. For example, the peak current **727** to a first memory array is delayed by a time **710a** from the peak current draw **729** for a second memory array. As a result, the aggregate current **730** peaks at roughly the same peaks of the individual current profiles **720b**, **722b**, and **724b** instead of being a sum of the peak currents as in FIG. 4B when refresh cycles overlap in time.

[0042] While the present disclosure contains many representative examples, these should not be construed as limitations on the scope of any disclosed methods or of what can be claimed, but rather as descriptions of features that may be specific to embodiments of disclosed methods. Certain features that are described in this document in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable sub-combination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination can be directed to a sub-combination or variation of a sub-combination.

[0043] Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. Moreover, the separation of various system components in the embodiments described in this patent document and attached appendices should not be understood as requiring such separation in all embodiments.

[0044] From the foregoing, it will be appreciated that specific embodiments of the present technology have been described herein for purposes of illustration, but that various modifications can be made without deviating from the scope of the technology. Further, while advantages associated with certain embodiments of the technology have been described in the context of those embodiments, other embodiments can also exhibit such advantages, and not all embodiments need necessarily exhibit such advantages to fall within the scope of the present technology. Accordingly, the present disclosure and associated technology can encompass other embodiments not expressly shown or described herein.

Claims

1. A method at a memory device, comprising: performing, at a first rank of the memory device, a first refresh operation; and performing, at a second rank of the memory device, a second refresh operation after a duration, wherein the duration is a product of a value and a refresh cycle time associated with refresh operations at the memory device.
2. The method of claim 1, further comprising: receiving a first refresh command directed to the first rank, wherein performing the first refresh operation is in accordance with receiving the first refresh

command; and receiving, after the duration from receiving the first refresh command, a second refresh command directed to the second rank, wherein performing the second refresh operation is in accordance with receiving the second refresh command.

3. The method of claim 2, wherein the duration is a predetermined quantity of time intervals between the first refresh command and the second refresh command.

4. The method of claim 2, wherein the duration is a minimum delay between the first refresh command and the second refresh command.

5. The method of claim 2, wherein performing the second refresh operation after the duration limits a maximum refresh current of the memory device, and wherein the maximum refresh current is a refresh current (IDD5) of the memory device.

6. The method of claim 1, wherein the first rank comprises a first logical rank of the memory device, and wherein the second rank comprises a second logical rank of the memory device.

7. The method of claim 1, wherein the memory device is a three-dimensional stacked (3DS) memory device.

8. A memory device, comprising: one or more ranks; and one or more controllers coupled with the one or more ranks and configured to cause the memory device to: perform, at a first rank of the one or more ranks, a first refresh operation; and perform, at a second rank of the one or more ranks, a second refresh operation after a duration, wherein the duration is a product of a value and a refresh cycle time associated with refresh operations at the memory device.

9. The memory device of claim 8, wherein the one or more controllers are configured to cause the memory device to: receive a first refresh command directed to the first rank, wherein performing the first refresh operation is in accordance with receiving the first refresh command; and receive, after the duration from receiving the first refresh command, a second refresh command directed to the second rank, wherein performing the second refresh operation is in accordance with receiving the second refresh command.

10. The memory device of claim 9, wherein the duration is a predetermined quantity of time intervals between the first refresh command and the second refresh command.

11. The memory device of claim 9, wherein the duration is a minimum delay between the first refresh command and the second refresh command.

12. The memory device of claim 9, wherein performing the second refresh operation after the duration limits a maximum refresh current of the memory device, and wherein the maximum refresh current is a refresh current (IDD5) of the memory device.

13. The memory device of claim 8, wherein the first rank comprises a first logical rank of the memory device, and wherein the second rank comprises a second logical rank of the memory device.

14. The memory device of claim 8, wherein the memory device is a three-dimensional stacked (3DS) memory device.

15. A method at a three-dimensional (3DS) memory device, comprising: performing, at a first rank of the 3DS memory device, a first refresh operation; and performing, at a second rank of the 3DS memory device, a second refresh operation after a duration, wherein the duration is a product of a value and a refresh cycle time associated with refresh operations at the 3DS memory device, and wherein performing the second refresh operation after the duration limits a maximum refresh current of the 3DS memory device.

16. The method of claim 15, further comprising: receiving a first refresh command directed to the first rank, wherein performing the first refresh operation is in accordance with receiving the first refresh command; and receiving, after the duration from receiving the first refresh command, a second refresh command directed to the second rank, wherein performing the second refresh operation is in accordance with receiving the second refresh command.

17. The method of claim 16, wherein the duration is a predetermined quantity of time intervals between the first refresh command and the second refresh command.

18. The method of claim 16, wherein the duration is a minimum delay between the first refresh command and the second refresh command.

19. The method of claim 16, wherein the maximum refresh current is a refresh current (IDD5) of the 3DS memory device.

20. The method of claim 15, wherein the first rank comprises a first logical rank of the 3DS memory device, and wherein the second rank comprises a second logical rank of the 3DS memory device.
