

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250266087

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Ramesh; Vijay S. et al.

EXTENDED MEMORY COMMUNICATION

Abstract

Systems, apparatuses, and methods related to extended memory communication subsystems for performing extended memory operations are described. An example method can include receiving, at a processing unit that is coupled between a host device and a non-volatile memory device, signaling indicative of a plurality of operations to be performed on data written to or read from the non-volatile memory device. The method can further include performing, at the processing unit, at least one operation of the plurality of operations in response to the signaling. The method can further include accessing a portion of a memory array in the non-volatile memory device. The method can further include transmitting additional signaling indicative of a command to perform one or more additional operations of the plurality of operations on the data written to or read from the non-volatile memory device.

Inventors: Ramesh; Vijay S. (Boise, ID), Porterfield; Allan (Durham, NC)

Applicant: Lodestar Licensing Group LLC (Evanston, IL)

Family ID: 1000008586874

Appl. No.: 19/202053

Filed: May 08, 2025

Related U.S. Application Data

parent US continuation 18478453 20230929 parent-grant-document US 12300315 child US 19202053

parent US continuation 17453136 20211101 parent-grant-document US 11810618 child US 18478453

parent US continuation 17109999 20201202 parent-grant-document US 11164625 child US 17453136

parent US continuation 16744541 20200116 parent-grant-document US 10910048 child US 17109999

Publication Classification

Int. Cl.: **G11C13/00** (20060101); **G06F7/57** (20060101); **G06F9/30** (20180101); **G06F9/54** (20060101)

U.S. Cl.:

CPC **G11C13/0023** (20130101); **G06F7/57** (20130101); **G06F9/3001** (20130101); **G06F9/542** (20130101); **G06F9/546** (20130101); **G11C13/004** (20130101); **G11C13/0069** (20130101); G11C2213/71 (20130101)

Background/Summary

PRIORITY INFORMATION [0001] This application is a continuation of U.S. application Ser. No. 18/478,453, filed Sep. 29, 2023, which is a continuation of U.S. application Ser. No. 17/453,136, filed Nov. 1, 2021, which issued as U.S. Pat. No. 11,810,618 on Nov. 7, 2023, which is a continuation of U.S. application Ser. No. 17/109,999, filed Dec. 2, 2020, which issued as U.S. Pat. No. 11,164,625 on Nov. 2, 2021, which is a continuation of U.S. application Ser. No. 16/744,541, filed Jan. 16, 2020, which issued as U.S. Pat. No. 10,910,048, issued Feb. 2, 2021, the contents of which are incorporated herein by reference.

TECHNICAL FIELD

[0002] The present disclosure relates generally to semiconductor memory and methods, and more particularly, to apparatuses, systems, and methods for an extended memory communication.

BACKGROUND

[0003] Memory devices are typically provided as internal, semiconductor, integrated circuits in computers or other electronic systems. There are many different types of memory including volatile and non-volatile memory. Volatile memory can require power to maintain its data (e.g., host data, error data, etc.) and includes random access memory (RAM), dynamic random access memory (DRAM), static random access memory (SRAM), synchronous dynamic random access memory (SDRAM), and thyristor random access memory (TRAM), among others. Non-volatile memory can provide persistent data by retaining stored data when not powered and can include NAND flash memory, NOR flash memory, and resistance variable memory such as phase change random access memory (PCRAM), resistive random access memory (RRAM), and magnetoresistive random access memory (MRAM), such as spin torque transfer random access memory (STT RAM), among others.

[0004] Memory devices may be coupled to a host (e.g., a host computing device) to store data, commands, and/or instructions for use by the host while the computer or electronic system is operating. For example, data, commands, and/or instructions can be transferred between the host and the memory device(s) during operation of a computing or other electronic system.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0005] FIG. 1 is a functional block diagram in the form of a computing system including an apparatus including a first plurality of communication subsystems, a second plurality of communication subsystems, and a plurality of memory devices in accordance with a number of embodiments of the present disclosure.

[0006] FIG. 2 is yet another functional block diagram in the form of a computing system including an apparatus including a first plurality of communication subsystems, a second plurality of communication subsystems, and a plurality of memory devices in accordance with a number of embodiments of the present disclosure.

[0007] FIG. 3 is yet another functional block diagram in the form of a computing system including an apparatus including a first plurality of communication subsystems, a second plurality of communication subsystems, and a plurality of memory devices in accordance with a number of embodiments of the present disclosure.

[0008] FIG. 4 is yet another functional block diagram in the form of a computing system including an apparatus including a first plurality of communication subsystems, a second plurality of communication subsystems, and a plurality of memory devices in accordance with a number of embodiments of the present disclosure.

[0009] FIG. 5 is a functional block diagram in the form of an apparatus of a computing core including a number of ports in accordance with a number of embodiments of the present disclosure.

[0010] FIG. 6 is a flow diagram representing an example method corresponding to extended memory communication in accordance with a number of embodiments of the present disclosure.

DETAILED DESCRIPTION

[0011] Systems, apparatuses, and methods related to extended memory communication subsystems for performing extended memory operations are described. An example method can include receiving, at a processing unit that is coupled between a host device and a non-volatile memory device, signaling indicative of a plurality of operations to be performed on data written to or read from the non-volatile memory device. The method can further include performing, at the processing unit, at least one operation of the plurality of operations in response to the signaling. The method can further include accessing a portion of a memory array in the non-volatile memory device. The method can further include transmitting additional signaling indicative of a command to perform one or more additional operations of the plurality of operations on the data written to or read from the non-volatile memory device.

[0012] Extended memory communication can include providing signals and/or commands across extended memory. An extended memory interface can transfer instructions to perform operations specified by a single address and operand and may be performed by the computing device that includes a processing unit and a memory resource. The computing device can perform extended memory operations on data streamed through the computing device without receipt of intervening commands. The extended memory operations can include an operation in which data is ordered, reordered, removed, or discarded, a comma-separated value parsing operation, or both. In an example, a computing device is configured to receive a command to perform an operation that comprises performing an operation on data with the processing unit of the computing device and determine that an operand corresponding to the operation is stored in the memory resource. The computing device can further perform the operation using the operand stored in the memory resource.

[0013] The computing device can perform hardware acceleration by sending instructions and/or commands to a number of hardware accelerators to perform the operation. In some examples, a portion of the operation can be sent to a first hardware accelerator and a second portion of the operation can be sent to a second hardware accelerator. In some examples, the operation can be sent to a hardware accelerator for completion and the hardware accelerator can send a portion of the operation to an additional hardware accelerator to complete a portion of the operation. In this way, results from more than one hardware accelerator can be sent to the computing device to combine the results or a primary hardware accelerator can combine the results and send the completed result to the computing device.

[0014] Hardware acceleration can be implemented in computing systems to perform certain tasks and/or functions in a manner that is more efficient (e.g., faster, more accurate, higher quality, etc.)

in comparison to performing the task and/or function using a central processing unit (CPU) of the computing system. For example, by providing dedicated hardware (e.g., a hardware accelerator or hardware acceleration unit) that is configured to perform a certain task and/or function that can otherwise be performed using the CPU of the computing system, certain tasks and/or functions can be processed in a more efficient manner than in approaches in which the CPU is responsible for performance of such tasks and/or functions. This can further allow for processing resources that could otherwise be consumed by the CPU to be freed up, thereby further improving performance of the computing system.

[0015] Some examples of hardware accelerators include sounds processing units (e.g., sound cards), graphics processing units (GPUs or “graphics cards”), digital signal processing units, analog signal processing units, computer networking processing units (e.g., networks on a chip, TCP offload engines, I/O acceleration processing units, etc.), cryptography processing units (e.g., cryptographic accelerator units, which can provide hardware-based encryption and/or decryption), artificial intelligence processing units (e.g., vision processing units, neural network processing units, etc.), tensor processing units, physics processing units, regular expression processing units, and/or data compression acceleration units, among others. Hardware accelerators can be provided as computer hardware in the form of a field-programmable gate array (FPGA), an application-specific integrated circuit (ASIC), and/or a complex programmable logic device, a system-on-chip, among others. It will be appreciated that the foregoing enumerated examples of hardware accelerators and specifically enumerated examples of computer hardware are neither limiting nor exhaustive, and other hardware accelerators and/or computer hardware are contemplated within the scope of the disclosure.

[0016] In some approaches, hardware accelerators can be deployed in a computing system as discrete components that perform a specified task and/or function with no visibility to other hardware accelerators that can be deployed within the computing system. For example, in some approaches, a hardware accelerator can operate without knowledge of other hardware accelerators deployed within the computing system. Further, in some approaches, hardware accelerators can be dedicated to perform a limited set of specific tasks and/or functions. For example, a sound processing unit can be provided in a computing system with the sole purpose of performing hardware acceleration on signals related to auditory playback for the computing system. As another example, a GPU can be provided in a computing system for the sole purpose of performing hardware acceleration on signals related to visual display for the computing system.

[0017] As is described below, the computing device can be a RISC-V application processor core, capable of supporting full-featured operating systems such as Linux. This particular core can be used in association with applications such as internet-of-things (IoT) nodes and gateways, storage, and/or networking. The core can be coupled to a number of ports, such as a memory port, a system port, a peripheral port, and/or a front port. As an example, the memory port can be in communication with a memory device, the system port can be in communication with an on-chip accelerator, the peripheral port can be in communication with an off-chip serial port, and/or the front port can be in communication with a host interface, as will be described further below in association with FIG. 4.

[0018] In this way, the first communication subsystems can be used to direct data from a particular port (e.g., a memory port of a computing device) through a first communication subsystem (e.g., a multiplexer that selects that particular memory port) and transfer it through a second communication subsystem (e.g., an interface such as an AXI interconnect interface) to a memory controller that transfer the data to a memory device (e.g., a DDR memory, a three-dimensional (3-D) cross-point memory, a NAND memory, etc.). In an example, the AXI interconnect interfaces can conform to the AMBA® AXI version 4 specifications from ARM®, including the AXI4-Lite control register interface subset.

[0019] As used herein, an “extended memory operation” refers to a memory operation that can be

specified by a single address (e.g., a memory address) and an operand, such as a 64-bit operand. An operand can be represented as a plurality of bits (e.g., a bit string or string of bits). Embodiments are not limited to operations specified by a 64-bit operand, however, and the operation can be specified by an operand that is larger (e.g., 128-bits, etc.) or smaller (e.g., 32-bits) than 64-bits. As described herein, the effective address space accessible with which to perform extended memory operations is the size of a memory device or file system accessible to a host computing system or storage controller.

[0020] Extended memory operations can include instructions and/or operations that can be performed by a processing device (e.g., by a processing device such as a core **110**, **210**, **310**, **410**, or a core computing device specifically shown as **510** in FIG. 5). Examples of a core can include a reduced instruction set computing device or other hardware processing device that can execute instructions to perform various computing tasks. In some embodiments, performing an extended memory operation can include retrieving data and/or instructions stored in a memory resource of the computing device, performing the operation within the computing device **110** (e.g., without transferring the data or instructions to circuitry external to the computing device), and storing the result of the extended memory operation in the memory resource of the computing device **110** or in secondary storage (e.g., in a memory device such as the memory device **116-1**, **116-2**, illustrated in FIG. 1, herein). Signaling indicative of a plurality of operations to be performed on data written to or from a memory device can be sent to or from the computing devices **110**, accelerators **114**, etc.

[0021] Non-limiting examples of extended memory operations can include floating point add accumulate, 32-bit complex operations, square root address (SQRT(addr)) operations, conversion operations (e.g., converting between floating-point and integer formats, and/or converting between floating-point and universal number formats such as Type I, Type II, and/or Type III universal number formats, posit formats, etc.), normalizing data to a fixed format, absolute value operations, etc. In some embodiments, extended memory operations can include operations performed by the computing device that update in place (e.g., in which a result of an extended memory operation is stored at the address in which an operand used in performance of the extended memory operation is stored prior to performance of the extended memory operation), as well as operations in which previously stored data is used to determine a new data (e.g., operations in which an operand stored at a particular address is used to generate new data that overwrites the particular address where the operand was stored).

[0022] As a result, in some embodiments, performance of extended memory operations can mitigate or eliminate locking or mutex operations, because the extended memory operation(s) can be performed within the computing device, which can reduce contention between multiple thread execution. Reducing or eliminating performance of locking or mutex operations on threads during performance of the extended memory operations can lead to increased performance of a computing system, for example, because extended memory operations can be performed in parallel within a same computing device or across two or more of the computing devices that are in communication with each other. In addition, in some embodiments, extended memory operations described herein can mitigate or eliminate locking or mutex operations when a result of the extended memory operation is transferred from the computing device that performed the operation to a host.

[0023] Memory devices may be used to store important or critical data in a computing device and can transfer, via at least one extended memory interface, such data between a host associated with the computing device. However, as the size and quantity of data stored by memory devices increases, transferring the data to and from the host can become time consuming and resource intensive. For example, when a host requests performance of memory operations using large blocks of data, an amount of time and/or an amount of resources consumed in obliging the request can increase in proportion to the size and/or quantity of data associated with the blocks of data.

[0024] As storage capability of memory devices increases, these effects can become more pronounced as more and more data are able to be stored by the memory device and are therefore

available for use in memory operations. In addition, because data may be processed (e.g., memory operations may be performed on the data), as the amount of data that is able to be stored in memory devices increases, the amount of data that may be processed can also increase. This can lead to increased processing time and/or increased processing resource consumption, which can be compounded in performance of certain types of memory operations. In order to alleviate these and other issues, embodiments herein can allow for extended memory operations to be performed using a memory device, one or more computing devices, and/or memory array(s) and a first plurality of communication subsystems (e.g., multiplexers) and a second plurality of subsystems (e.g., interfaces such as AXI interconnects) in order to transfer data more efficiently from a computing device to a memory device and/or from a computing device to a host, and vice versa.

[0025] In some approaches, performing memory operations can require multiple clock cycles and/or multiple function calls to memory of a computing system such as a memory device and/or memory array. In contrast, embodiments herein can allow for performance of extended memory operations in which a memory operation is performed with a single function call or command. For example, in contrast to approaches in which at least one command and/or function call is utilized to load data to be operated upon and then at least one subsequent function call or command to store the data that has been operated upon is utilized, embodiments herein can allow for performance of memory operations using fewer function calls or commands in comparison to other approaches. Further, the computing devices of the computing system can receive requests to perform the memory operations via a first communication subsystem (e.g., a multiplexer, a control network-on-chip, etc.) and/or a second communication subsystem (e.g., an interface, an interconnect such as an AXI interconnect, etc.) and can receive blocks of data for executing the requested memory operations from the memory device via the first communication subsystem and the second communication subsystem. While the first and the second communication subsystem are described in tandem, embodiments are not so limited. As an example, the requests for data and/or receipt of blocks of data can be via the second communication subsystem alone.

[0026] By reducing the number of function calls and/or commands utilized in performance of memory operations, an amount of time consumed in performing such operations and/or an amount of computing resources consumed in performance of such operations can be reduced in comparison to approaches in which multiple function calls and/or commands are required for performance of memory operations. Further, embodiments herein can reduce movement of data within a memory device and/or memory array because data may not need to be loaded into a specific location prior to performance of memory operations. This can reduce processing time in comparison to some approaches, especially in scenarios in which a large amount of data is subject to a memory operation.

[0027] Further, extended memory operations described herein can allow for a much larger set of type fields in comparison to some approaches. For example, an instruction executed by a host to request performance of an operation using data in a memory device (e.g., a memory sub-system) can include a type, an address, and a data field. The instruction can be sent to at least one of a plurality of computing devices via a first communication subsystem (e.g., a multiplexer) and a second communication subsystem (e.g., an interface) and the data can be transferred from the memory device via the first and/or second communication subsystem. The type field can correspond to the particular operation being requested, the address can correspond to an address in which data to be used in performance of the operation is stored, and the data field can correspond to the data (e.g., an operand) to be used in performing the operation. In some approaches, type fields can be limited to different size reads and/or writes, as well as some simple integer accumulate operations. In contrast, embodiments herein can allow for a broader spectrum of type fields to be utilized because the effective address space that can be used when performing extended memory operations can correspond to a size of the memory device. By extending the address space available to perform operations, embodiments herein can therefore allow for a broader range of

type fields and, therefore, a broader spectrum of memory operations can be performed than in approaches that do not allow for an effective address space that is the size of the memory device. [0028] In the following detailed description of the present disclosure, reference is made to the accompanying drawings that form a part hereof, and in which is shown by way of illustration how one or more embodiments of the disclosure may be practiced. These embodiments are described in sufficient detail to enable those of ordinary skill in the art to practice the embodiments of this disclosure, and it is to be understood that other embodiments may be utilized and that process, electrical, and structural changes may be made without departing from the scope of the present disclosure.

[0029] As used herein, designators such as “X,” “Y,” “N,” “M,” “A,” “B,” “C,” “D,” etc., particularly with respect to reference numerals in the drawings, indicate that a number of the particular feature so designated can be included. It is also to be understood that the terminology used herein is for the purpose of describing particular embodiments only, and is not intended to be limiting. As used herein, the singular forms “a,” “an,” and “the” can include both singular and plural referents, unless the context clearly dictates otherwise. In addition, “a number of,” “at least one,” and “one or more” (e.g., a number of memory banks) can refer to one or more memory banks, whereas a “plurality of” is intended to refer to more than one of such things. Furthermore, the words “can” and “may” are used throughout this application in a permissive sense (i.e., having the potential to, being able to), not in a mandatory sense (i.e., must). The term “include,” and derivations thereof, means “including, but not limited to.” The terms “coupled” and “coupling” mean to be directly or indirectly connected physically or for access to and movement (transmission) of commands and/or data, as appropriate to the context. The terms “data” and “data values” are used interchangeably herein and can have the same meaning, as appropriate to the context.

[0030] The figures herein follow a numbering convention in which the first digit or digits correspond to the figure number and the remaining digits identify an element or component in the figure. Similar elements or components between different figures may be identified by the use of similar digits. For example, **104** may reference element “04” in FIG. **1**, and a similar element may be referenced as **204** in FIG. **2**. A group or plurality of similar elements or components may generally be referred to herein with a single element number. For example, a plurality of reference elements **110-1**, **110-2**, **110-3**, **110-4**, **110-5** may be referred to generally as **110**. As will be appreciated, elements shown in the various embodiments herein can be added, exchanged, and/or eliminated so as to provide a number of additional embodiments of the present disclosure. In addition, the proportion and/or the relative scale of the elements provided in the figures are intended to illustrate certain embodiments of the present disclosure and should not be taken in a limiting sense.

[0031] FIG. **1** is a functional block diagram in the form of a computing system **100** including an apparatus **104** including a plurality of computing devices **110**, a first plurality of communication subsystems **108**, a second plurality of communication subsystems **106**, a plurality of hardware accelerators **114**, and a plurality of memory devices **116**, in accordance with a number of embodiments of the present disclosure. As used herein, an “apparatus” can refer to, but is not limited to, any of a variety of structures or combinations of structures, such as a circuit or circuitry, a die or dice, a module or modules, a device or devices, or a system or systems, for example. In the embodiment illustrated in FIG. **1**, memory devices **116-1** . . . **116-N** can include one or more memory modules (e.g., double data rate (DDR) memory, three-dimensional (3D) cross-point memory, NAND memory, single in-line memory modules, dual in-line memory modules, etc.). The memory devices **116-1**, . . . , **116-N** can include volatile memory and/or non-volatile memory. In a number of embodiments, memory devices **116-1**, . . . , **116-N** can include a multi-chip device. A multi-chip device can include a number of different memory types and/or memory modules. For example, a memory system can include non-volatile or volatile memory on any type of a module.

[0032] The memory devices **116-1**, . . . , **116-N** can provide main memory for the computing system **100** or could be used as additional memory or storage throughout the computing system **100**. Each memory device **116-1**, . . . , **116-N** can include one or more arrays of memory cells, e.g., volatile and/or non-volatile memory cells. The arrays can be flash arrays with a NAND architecture, for example. Embodiments are not limited to a particular type of memory device. For instance, the memory device can include RAM, ROM, DRAM, SDRAM, PCRAM, RRAM, and flash memory, among others.

[0033] In embodiments in which the memory devices **116-1**, . . . , **116-N** include non-volatile memory, the memory devices **116-1**, . . . , **116-N** can be flash memory devices such as NAND or NOR flash memory devices. Embodiments are not so limited, however, and the memory devices **116-1**, . . . , **116-N** can include other non-volatile memory devices such as non-volatile random-access memory devices (e.g., NVRAM, ReRAM, FeRAM, MRAM, PCM), “emerging” memory devices such as 3-D Crosspoint (3D XP) memory devices, etc., or combinations thereof. A 3D XP array of non-volatile memory can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many flash-based memories, 3D XP non-volatile memory can perform a write in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased.

[0034] As illustrated in FIG. 1, the plurality of computing devices **110-1**, **110-2**, **110-3**, **110-4**, **110-5** (hereinafter referred to collectively as plurality of computing devices **110**) can be coupled to an SRAM **109**. The plurality of computing devices **110** can be coupled to the SRAM **109** through a bus matrix. Further, the plurality of computing devices **110** can be coupled to the first plurality of communication subsystems (e.g., multiplexers) **108-1**, **108-2**. The first plurality of communication subsystems **108** can include circuitry and/or logic configured to allocate and de-allocate resources to the computing devices **110** during performance of operations described herein. For example, the circuitry and/or logic can allocate and/or de-allocate resources to the computing devices **110** during performance of extended memory operations described herein.

[0035] The plurality of computing devices **110** can be coupled to a first (**108-1**) of the first plurality of communication subsystems **108** through the SRAM **109**. The plurality of computing devices **110** can be directly coupled to the first of the first plurality of communication subsystems **108-1** and/or to a second (**108-2**) of the first plurality of communication subsystems **108**, as illustrated by arrows in FIG. 1. In this way, each of the first plurality of communication subsystems can select a particular computing device **110** for transferring data, and vice versa, each of the computing devices **110** can transfer data through the first plurality of communication subsystems **108**.

[0036] The first plurality of communication subsystems **108-1** can be coupled to a second plurality of communication subsystems (e.g., interfaces such as an interconnect interface) **106-1**, **106-2**, **106-3**, **106-4**, **106-5** (hereinafter referred to collectively as second plurality of communication subsystems **106**). Each of the second plurality of communication subsystems **106** can be coupled to a corresponding one of a controller **112**, an accelerator **114**, and a host interface **120**. In one example, the second plurality of communication subsystems **106** can be coupled to the corresponding controller **112**, accelerators **114**, and/or host interface **120** via a number of AXI buses.

[0037] As is illustrated, a first (**106-1**) of the second plurality of communication subsystems **106** can be coupled to the controller (e.g., memory controller) **112**. The controller **112** can be coupled to a number of memory devices **116-1**, . . . , **116-N** via a number of channels **107-1**, . . . , **107-N**. A second (**106-2**), third (**106-3**), and fourth (**106-4**) of the second plurality of communication subsystems **106** can each be coupled to a corresponding one of the plurality of hardware accelerators **114-1**, **114-2**, **114-3**. The communication subsystem **108-1** can be coupled to the second plurality of communication subsystems **106-2**, **106-3**, **106-4** via respective buffers **119-1**, **119-2**, **119-3**. The second plurality of communication subsystems **106-2**, **106-3**, **106-4** can be

coupled to the plurality of hardware accelerators **114** via respective buffers **117-1**, **117-2**, **117-3**. The hardware accelerators **114** can be used for performing a number of posit operations, and/or for communication with an internal SRAM on the FPGA.

[0038] A posit operation can refer to an operation performed using universal number (“unum”) formatted bit strings as operands and/or as inputs. As used herein, universal number formatted bit strings can provide an alternative to the IEEE floating point bit string standard. Several universal number formats exist (e.g., Type I universal numbers, Type II universal numbers, and Type III universal numbers). The Type III unum format is referred to herein as a “posit format” or, for simplicity, a “posit.” In contrast to floating-point bit strings, posits can, under certain conditions, allow for higher precision (e.g., a broader dynamic range, higher resolution, and/or higher accuracy) than floating-point numbers with the same bit width. This can allow for operations performed by a computing system to be performed at a higher rate (e.g., faster) when using posits than with floating-point numbers, which, in turn, can improve the performance of the computing system by, for example, reducing a number of clock cycles used in performing operations thereby reducing processing time and/or power consumed in performing such operations. In addition, the use of posits in computing systems can allow for higher accuracy and/or precision in computations than floating-point numbers, which can further improve the functioning of a computing system in comparison to some approaches (e.g., approaches which rely upon floating-point format bit strings).

[0039] Posits can be highly variable in precision and accuracy based on the total quantity of bits and/or the quantity of sets of integers or sets of bits included in the posit. In addition, posits can generate a wide dynamic range. The accuracy, precision, and/or the dynamic range of a posit can be greater than that of a float, or other numerical formats, under certain conditions, as described in more detail herein. The variable accuracy, precision, and/or dynamic range of a posit can be manipulated, for example, based on an application in which a posit will be used. In addition, posits can reduce or eliminate the overflow, underflow, NaN, and/or other corner cases that are associated with floats and other numerical formats. Further, the use of posits can allow for a numerical value (e.g., a number) to be represented using fewer bits in comparison to floats or other numerical formats.

[0040] A computing device **110** can send a command to perform a posit operation and/or additional operations. The computing device **100** can divide the posit operation into portions of sub-operations to each be sent to a hardware accelerator **114**. For example, a first computing device **110-1** can divide a posit operation into two sub-operations and a first of the two sub-operations can be sent to a first hardware accelerator **114-1** and a second of the two sub-operations can be sent to a second hardware accelerator **114-2**. The results of the first and the second sub-operation can be sent to the first computing device **110-1** and the first computing device **110-1** can combine the results into a single result of the posit operation.

[0041] In one example, a computing device **110** can send a posit operation to a first hardware accelerator **114-1** and the first hardware accelerator **114-1** can send a portion of the posit operation to a second hardware accelerator **114-2**. Upon receipt of the result of the portion of the posit operation from the second hardware accelerator **114-2**, the first hardware accelerator **114-1** can generate a result for the posit operation, including the result from the portion of the posit operation. Likewise, any number of divisions of a posit operation can be sent from a computing device **110** to particular numbers of corresponding hardware accelerators **114** to perform the posit operation and the results can be combined for a final result of the posit operation. Likewise, multiple computing devices **110** can sub-divide a posit operation and send different portions to different hardware accelerators **114** to perform the sub-divided posit operations. As an example, a posit operation can be sub-divided by computing devices **110** and the sub-divided posit operations can be further sub-divided by each corresponding computing device **110** and sent to different hardware accelerators **114**.

[0042] Further, another additional one (also not illustrated) of the second plurality of communication subsystems **106** can be used for transferring data off-chip through an off-chip serial port. The fifth (**106-5**) of the second plurality of communication subsystems **106** can be coupled to a host interface **120** and can communicate, via channels **103/105**, with a host controller **101** of a host **102**. While not illustrated, a communication subsystem (such as another of the second plurality of communication subsystems, not illustrated) can be coupled to logic circuitry. The logic circuitry can be on a same field programmable gate array (FPGA) as the computing devices **110**, first plurality of communication subsystems, second plurality of communication subsystems **106**, etc.

[0043] In one embodiment, the computing device **110** can process an operation queue of messages from a host **102**. The operation queue can be processed by the computing device **110** by reading an input message with arguments and executing a desired function. Further, the computing device **110** can read and/or write data to at least one of the memory devices **116-1**, . . . , **116-N** in order to perform an operation. The computing device **110** can generate a work item message to be performed and generate a message to send to at least one of the hardware accelerators **114** indicating to perform a work item associated with the input message. The message can identify an operation or sub-operation to be performed in relation to the receipted input message and arguments, identify a hardware accelerator to be activated and a function to be performed, identify an input data location of the memory device **116**, and identify an output data location of the memory device **116**. As an example, the input data location can indicate a location of data in the memory device **116** to retrieve data from in order to perform the work item. The output data location can indicate a location in the memory device **116** to store the resultant output of the operation of the work item. The work item message can be sent to the corresponding hardware accelerator **114** queue (e.g., such as to the buffer **117** of a respective hardware accelerator **114**). As the results of the operation are generated or received, additional messages indicating additional operations to be performed can be generated and sent to hardware accelerators. The generation of messages and reception of results can continue until a final result of the initial operation brings the work item to completion. Upon completion of the work item, a completion message can be sent to the computing device **110** indicating the work item has been complete. A message can be sent to the host **102** indicating that the work item has been completed.

[0044] The hardware accelerator **114**, upon receipt of a work item, can read the work item message, including corresponding data locations in the memory device **116**. The hardware accelerator **114** can perform the requested accelerator operations contained within the message. In one embodiment, the hardware accelerator **114** can send a portion of the operation to an additional hardware accelerator **114** (e.g., hardware accelerator **114-1** can receive the message and can send a portion of the operation in the message to hardware accelerator **114-2** to be completed by hardware accelerator **114-2**). The completed portion of the operation (executed by hardware accelerator **114-2**) can be sent to the initial hardware accelerator (**114-1**) and the initial hardware accelerator (**114-1**) can combine the completed portion with other results to finalized completion of the operation in the work item message. Once fully complete, the hardware accelerator **114** can send a message to the computing device **110** indicating the work item has been completed.

[0045] In one embodiment, the host **102** can send a request for a computing device **110** to perform an operation. The computing device **110** can perform data initialization and write the data to location in the memory device **116**. As the computing device **110** generates 4K (or a multiple thereof) of data, the computing device **110** can create a work item to be completed by a hardware accelerator **114**. To further process the data. When the hardware accelerator **114** completes the work item, the hardware accelerator can send a message to the computing device **110** that the work item is complete. The computing device **110** can either further process the data, send the data to another hardware accelerator **114**, or leave the data in the memory device **116** alone and continue processing the data further.

[0046] The host **102** can map the data into its own address space. The host **102** can map a file (e.g., Linux for example) into the computing device **110** processing unit address space. The computing device **110** has a map between its addresses and the locations within the memory device **116**. When a hardware accelerator **114** work item is created, the address passed to the hardware accelerator **114** can be the logical block address of the memory device **116**. The host **102** can be responsible for mapping the address between the file system and the 64-bit address space of the computing device **110**. The computing device **110** can be responsible for mapping its addresses into logical block locations of the memory device **116**. In this way, the hardware accelerators **114** are responsible for transferring data from one logical data location of the memory device **116** to another.

[0047] The host **102** can be a host system such as a personal laptop computer, a desktop computer, a digital camera, a smart phone, a memory card reader, and/or internet-of-things enabled device, among various other types of hosts, and can include a memory access device, e.g., a processor (or processing device). One of ordinary skill in the art will appreciate that “a processor” can intend one or more processors, such as a parallel processing system, a number of coprocessors, etc. The host **102** can include a system motherboard and/or backplane and can include a number of processing resources (e.g., one or more processors, microprocessors, or some other type of controlling circuitry). In some embodiments, the host can include the host controller **101**, which can be configured to control at least some operations of the host **102** by, for example, generating and transferring commands to the host controller to cause performance of operations such as extended memory operations. The host controller **101** can include circuitry (e.g., hardware) that can be configured to control at least some operations of the host **102**. For example, the host controller **101** can be an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), or other combination of circuitry and/or logic configured to control at least some operations of the host **102**.

[0048] The system **100** can include separate integrated circuits or the host **102**, the first plurality of communication subsystems **108**, the second plurality of communication subsystems **106**, the controller **112**, the on-chip accelerators **114**, the host interface **120**, and/or the memory devices **116-1**, . . . , **116-N** can be on the same integrated circuit. The system **100** can be, for instance, a server system and/or a high performance computing (HPC) system and/or a portion thereof. Although the example shown in FIG. 1 illustrate a system having a Von Neumann architecture, embodiments of the present disclosure can be implemented in non-Von Neumann architectures, which may not include one or more components (e.g., CPU, ALU, etc.) often associated with a Von Neumann architecture.

[0049] The controller **112** can be configured to request a block of data from one or more of the memory devices **116-1**, . . . , **116-N** and cause the cores **110-1**, . . . , **110-N**, which may be referred to in the alternative as “computing devices,” herein, to perform an operation (e.g., an extended memory operation) on the block of data. The operation may be performed to evaluate a function that can be specified by a single address and one or more operands associated with the block of data. The controller **112** can be further configured to cause a result of the extended memory operation to be stored in one or more of the computing devices **110-1**, . . . , **110-N** through the second **106** and/or the first **108** communication subsystems and/or to be transferred to a channel (e.g., communication paths **103** and/or **105**) and/or the host **102**.

[0050] In some embodiments, the second plurality of communication subsystems **106** can request a remote command, start a DMA command, send a read/write location, and/or send a start function execution command to one of the plurality of computing devices **110**. In some embodiments, the second plurality of communication subsystems **106** can request that a block of data be copied from a buffer of a computing device **110** to a buffer of a memory controller **112** or memory device **116**. Vice versa, one of the second plurality of communication subsystems **106** can request that a block of data be copied to the buffer of the computing device **110** from the buffer of the memory controller **112** or memory device **116**. The second plurality of communication subsystems **106** can

request that a block of data be copied to a computing device **110** from a buffer of the host **102** or, vice versa, request that a block of data be copied from a computing device **110** to a host **102**. The second plurality of communication subsystems **106** can request that a block of data be copied to a buffer of the host **102** from a buffer of the memory controller **112** or memory device **116**. Vice versa, the second plurality of communication subsystems **106** can request that a block of data be copied from a buffer of the host **102** to a buffer of the memory controller **112** or memory device **116**. Further, in some embodiments, the second plurality of communication subsystems **106** can request that a command from a host be executed on a computing device **110**. The second plurality of communication subsystems **106** can request that a command from a computing device **110** be executed on an additional computing device **110**. The second plurality of communication subsystems **106** can request that a command from a memory controller **112** be executed on a computing device **110**. In some embodiments, as described in more detail in connection with FIG. 3, herein, the second plurality of communication subsystems **106** can include at least a portion of a controller (not illustrated).

[0051] In some embodiments, the second plurality of communication subsystems **106** can transfer a block of data (e.g., a direct memory access (DMA) block of data) from a computing device **110** to a media device **116** (via the memory controller **112**) or, vice versa, can transfer a block of data to a computing device **110** from a media device **116**. The second plurality of communication subsystems **106** transfer a block of data (e.g., a DMA block) from a computing device **110** to a host **102** or, vice versa, to a computing device **110** from a host **102**. Further, the second plurality of communication subsystems **106** can transfer a block of data (e.g., a DMA block) from a host **102** to a media device **116** or, vice versa, to a host **102** from a media device **116**. In some embodiments, the second plurality of communication subsystems **106** can receive an output (e.g., data on which an extended memory operation has been performed) from the computing devices **110-1**, . . . , **110-N** and transfer the output from the computing devices **110-1**, . . . , **110-N** to a controller **115** of the apparatus **104** and/or the host **102**, and vice versa. For example, the second plurality of communication subsystems **106** may be configured to receive data that has been subjected to an extended memory operation by the computing devices **110-1**, . . . , **110-N** and transfer the data that corresponds to the result of the extended memory operation to a controller **115** and/or the host **102**. In some embodiments, second plurality of communication subsystems **106** can include at least a portion of the controller **115**. For example, the second plurality of communication subsystems **106** can include the circuitry that comprises the controller **115**, or a portion thereof.

[0052] The memory controller **112** can be a “standard” or “dumb” memory controller. For example, the memory controller **112** can be configured to perform simple operations such as copy, write, read, error correct, etc. for the memory devices **116-1**, . . . , **116-N**. However, in some embodiments, the memory controller **112** does not perform processing (e.g., operations to manipulate data) on data associated with the memory devices **116-1**, . . . , **116-N**. For example, the memory controller **112** can cause a read and/or write operation to be performed to read or write data from or to the memory devices **116-1**, . . . , **116-N** via the communication paths **107-1**, . . . , **107-N**, but the memory controller **112** may not perform processing on the data read from or written to the memory devices **116-1**, . . . , **116-N**. In some embodiments, the memory controller **112** can be a non-volatile memory controller, although embodiments are not so limited.

[0053] The embodiment of FIG. 1 can include additional circuitry that is not illustrated so as not to obscure embodiments of the present disclosure. For example, the apparatus **104** can include address circuitry to latch address signals provided over I/O connections through I/O circuitry. Address signals can be received and decoded by a row decoder and a column decoder to access the memory devices **116-1**, . . . , **116-N**. It will be appreciated by those skilled in the art that the number of address input connections can depend on the density and architecture of the memory devices **116-1**, . . . , **116-N**.

[0054] In some embodiments, extended memory operations can be performed using the computing

system **100** shown in FIG. **1** by selectively storing or mapping data (e.g., a file) into a computing device **110**. The data can be selectively stored in an address space of the computing memory. In some embodiments, the data can be selectively stored or mapped in the computing device **110** in response to a command received from the host **102**. In embodiments in which the command is received from the host **102**, the command can be transferred to the computing device **110** via an interface (e.g., communication paths **103** and/or **105**) associated with the host **102** and via the first and second plurality of communication subsystems **108** and **106**, respectively. The interface(s) **103/105**, first plurality of communication subsystems **108** and the second plurality of communication subsystems **106** can be peripheral component interconnect express (PCIe) buses, double data rate (DDR) interfaces, interconnect interfaces (such as AXI interconnect interfaces), multiplexers (muxes), or other suitable interfaces or buses. Embodiments are not so limited, however.

[0055] In a non-limiting example in which the data (e.g., in which data to be used in performance of an extended memory operation) is mapped into the computing device **110**, the host controller **101** can transfer a command to the computing device **110** to initiate performance of an extended memory operation using the data mapped into the computing device **110**. In some embodiments, the host controller **101** can look up an address (e.g., a physical address) corresponding to the data mapped into the computing device **110** and determine, based on the address, which computing device (e.g., the computing device **110-1**) the address (and hence, the data) is mapped to. The command can then be transferred to the computing device (e.g., the computing device **110-1**) that contains the address (and hence, the data).

[0056] In some embodiments, the data can be a 64-bit operand, although embodiments are not limited to operands having a specific size or length. In an embodiment in which the data is a 64-bit operand, once the host controller **101** transfers the command to initiate performance of the extended memory operation to the correct computing device (e.g., the computing device **110-1**) based on the address at which the data is stored, the computing device (e.g., the computing device **110-1**) can perform the extended memory operation using the data.

[0057] In some embodiments, the computing devices **110** can be separately addressable across a contiguous address space, which can facilitate performance of extended memory operations as described herein. That is, an address at which data is stored, or to which data is mapped, can be unique for all the computing devices **110** such that when the host controller **101** looks up the address, the address corresponds to a location in a particular computing device (e.g., the computing device **110-1**).

[0058] For example, a first computing device **110-1** can have a first set of addresses associated therewith, a second computing device **110-2** can have a second set of addresses associated therewith, a third computing device **110-3** can have a third set of addresses associated therewith, through the n-th computing device (e.g., the computing device **110-N**), which can have an n-th set of addresses associated therewith. That is, the first computing device **110-1** can have a set of addresses 0000000 to 0999999, the second computing device **110-2** can have a set of addresses 1000000 to 1999999, the third computing device **110-3** can have a set of addresses 2000000 to 2999999, etc. It will be appreciated that these address numbers are merely illustrative, non-limiting, and can be dependent on the architecture and/or size (e.g., storage capacity) of the computing devices **110**.

[0059] As a non-limiting example in which the extended memory operation comprises a floating-point-add-accumulate operation (FLOATINGPOINT_ADD_ACCUMULATE), the computing devices **110** can treat the destination address as a floating-point number, add the floating-point number to the argument stored at the address of the computing device **110**, and store the result back in the original address. For example, when the host controller **101** (or an apparatus controller **115**, not shown) initiates performance of a floating-point add accumulate extended memory operation, the address of the computing device **110** that the host looks up (e.g., the address in the computing

device to which the data is mapped) can be treated as a floating-point number and the data stored in the address can be treated as an operand for performance of the extended memory operation. Responsive to receipt of the command to initiate the extended memory operation, the computing device **110** to which the data (e.g., the operand in this example) is mapped can perform an addition operation to add the data to the address (e.g., the numerical value of the address) and store the result of the addition back in the original address of the computing device **110**.

[0060] As described above, performance of such extended memory operations can, in some embodiments require only a single command (e.g., request command) to be transferred from the host **102** (e.g., from the host controller **101**) to the memory device **104** or from the controller **115** to the computing device(s) **110**. In contrast to some previous approaches, this can reduce an amount of time, for example, for multiple commands to traverse the interface(s) **103**, **105** and/or for data, such as operands to be moved from one address to another within the computing device(s) **110**, consumed in performance of operations.

[0061] In addition, performance of extended memory operations in accordance with the disclosure can further reduce an amount of processing power or processing time since the data mapped into the computing device **110** in which the extended memory operation is performed can be utilized as an operand for the extended memory operation and/or the address to which the data is mapped can be used as an operand for the extended memory operation, in contrast to approaches in which the operands must be retrieved and loaded from different locations prior to performance of operations. That is, at least because embodiments herein allow for loading of the operand to be skipped, performance of the computing system **100** may be improved in comparison to approaches that load the operands and subsequently store a result of an operations performed between the operands.

[0062] Further, in some embodiments, because the extended memory operation can be performed within a computing device **110** using the address and the data stored in the address and, in some embodiments, because the result of the extended memory operation can be stored back in the original address, locking or mutex operations may be relaxed or not required during performance of the extended memory operation. Reducing or eliminating performance of locking or mutex operations on threads during performance of the extended memory operations can lead to increased performance of the computing system **100** because extended memory operations can be performed in parallel within a same computing device **110** or across two or more of the computing devices **110**.

[0063] In some embodiments, valid mappings of data in the computing devices **110** can include a base address, a segment size, and/or a length. The base address can correspond to an address in the computing device **110** in which the data mapping is stored. The segment size can correspond to an amount of data (e.g., in bytes) that the computing system **100** can process, and the length can correspond to a quantity of bits corresponding to the data. It is noted that, in some embodiments, the data stored in the computing device(s) **110** can be uncacheable on the host **102**. For example, the extended memory operations can be performed entirely within the computing devices **110** without encumbering or otherwise transferring the data to or from the host **102** during performance of the extended memory operations.

[0064] In a non-limiting example in which the base address is 4096, the segment size is 1024, and the length is 16,386, a mapped address, 7234, may be in a third segment, which can correspond to a third computing device (e.g., the computing device **110-3**) among the plurality of computing devices **110**. In this example, the host **102** and/or the first **108** and second **106** communication subsystems can forward a command (e.g., a request) to perform an extended memory operation to the third computing device **110-3**. The third computing device **110-3** can determine if data is stored in the mapped address in a memory of the third computing device **110-3**. If data is stored in the mapped address (e.g., the address in the third computing device **110-3**), the third computing device **110-3** can perform a requested extended memory operation using that data and can store a result of the extended memory operation back into the address in which the data was originally stored.

[0065] In some embodiments, the computing device **110** that contains the data that is requested for performance of an extended memory operation can be determined by the host controller **101**, and/or the first **108** and/or second **106** communication subsystems. For example, a portion of a total address space available to all the computing devices **110** can be allocated to each respective computing device. Accordingly, the host controller **101** and/or the first **108** and/or second **106** communication subsystems can be provided with information corresponding to which portions of the total address space correspond to which computing devices **110** and can therefore direct the relevant computing devices **110** to perform extended memory operations. In some embodiments, the host controller **101** and/or the second **106** communication subsystems can store addresses (or address ranges) that correspond to the respective computing devices **110** in a data structure, such as a table, and direct performance of the extended memory operations to the computing devices **110** based on the addresses stored in the data structure.

[0066] Embodiments are not so limited, however, and in some embodiments, the host controller **101** and/or the second communication subsystems **106** can determine a size (e.g., an amount of data) of the memory resource(s) and, based on the size of the memory resource(s) associated with each computing device **110** and the total address space available to all the computing devices **110**, determine which computing device **110** stores data to be used in performance of an extended memory operation. In embodiments in which the host controller **101** and/or the second communication subsystems **106** determine the computing device **110** that stores the data to be used in performance of an extended memory operation based on the total address space available to all the computing devices **110** and the amount of memory resource(s) available to each computing device **110**, it can be possible to perform extended memory operations across multiple non-overlapping portions of the computing device memory resource(s).

[0067] Continuing with the above example, if there is not data in the requested address, the third computing device **110-3** can request the data as described in more detail in connection with FIGS. 2-5, herein, and perform the extended memory operation once the data is loaded into the address of the third computing device **110-3**. In some embodiments, once the extended memory operation is completed by the computing device (e.g., the third computing device **110-3** in this example), and/or the host **102** can be notified and/or a result of the extended memory operation can be transferred to the memory devices **116** and/or the host **102**.

[0068] In some embodiments, the memory controller **112** can be configured to retrieve blocks of data from a memory device(s) **116-1**, . . . , **116-N** coupled to the apparatus **104** in response to a request from a controller of the apparatus **104** or a host **102**. The memory controller **112** can subsequently cause the blocks of data to be transferred to the computing devices **110-1**, . . . , **110-N** and/or the apparatus controller. Similarly, the memory controller **112** can be configured to receive blocks of data from the computing devices **110** and/or the controller **115**. The memory controller **112** can subsequently cause the blocks of data to be transferred to a memory device **116** coupled to the storage controller **104**.

[0069] The blocks of data can be approximately 4 kilobytes in size (although embodiments are not limited to this particular size) and can be processed in a streaming manner by the computing devices **110-1**, . . . , **110-N** in response to one or more commands generated by the controller **115** and/or a host and sent via the second communication subsystems **106**. In some embodiments, the blocks of data can be 32-bit, 64-bit, 128-bit, etc. words or chunks of data, and/or the blocks of data can correspond to operands to be used in performance of an extended memory operation.

[0070] For example, as described in more detail in connection with FIGS. 2-5, herein, because the computing devices **110** can perform an extended memory operation (e.g., process) a second block of data in response to completion of performance of an extended memory operation on a preceding block of data, the blocks of data can be continuously streamed through the computing devices **110** while the blocks of data are being processed by the computing devices **110**. In some embodiments, the blocks of data can be processed in a streaming fashion through the computing devices **110** in

the absence of an intervening command from the controller and/or the host **102**. That is, in some embodiments, the controller **115** (or host **102**) can issue a command to cause the computing devices **110** to process blocks of data received thereto and blocks of data that are subsequently received by the computing devices **110** can be processed in the absence of an additional command from the controller.

[0071] In some embodiments, processing the blocks of data can include performing an extended memory operation using the blocks of data. For example, the computing devices **110-1**, . . . , **110-N** can, in response to commands from the controller via the second plurality of communication subsystems **106**, perform extended memory operations the blocks of data to evaluate one or more functions, remove unwanted data, extract relevant data, or otherwise use the blocks of data in connection with performance of an extended memory operation.

[0072] In a non-limiting example in which the data (e.g., in which data to be used in performance of an extended memory operation) is mapped into one or more of the computing devices **110**, the controller can transfer a command to the computing device **110** to initiate performance of an extended memory operation using the data mapped into the computing device(s) **110**. In some embodiments, the controller **115** can look up an address (e.g., a physical address) corresponding to the data mapped into the computing device(s) **110** and determine, based on the address, which computing device (e.g., the computing device **110-1**) the address (and hence, the data) is mapped to. The command can then be transferred to the computing device (e.g., the computing device **110-1**) that contains the address (and hence, the data). In some embodiments, the command can be transferred to the computing device (e.g., the computing device **110-1**) via the second communication subsystem **106**.

[0073] The controller **115** (or a host) can be further configured to send commands to the computing devices **110** to allocate and/or de-allocate resources available to the computing devices **110** for use in performing extended memory operations using the blocks of data. In some embodiments, allocating and/or de-allocating resources available to the computing devices **110** can include selectively enabling some of the computing devices **110** while selectively disabling some of the computing devices **110**. For example, if less than a total number of computing devices **110** are required to process the blocks of data, the controller **115** can send a command to the computing devices **110** that are to be used for processing the blocks of data to enable only those computing devices **110** desired to process the blocks of data.

[0074] The controller **115** can, in some embodiments, be further configured to send commands to synchronize performance of operations, such as extended memory operations, performed by the computing devices **110**. For example, the controller **115** (and/or a host) can send a command to a first computing device **110-1** to cause the first computing device **110-1** to perform a first extended memory operation, and the controller **115** (or the host) can send a command to a second computing device **110-2** to perform a second extended memory operation using the second computing device. Synchronization of performance of operations, such as extended memory operations, performed by the computing devices **110** by the controller **115** can further include causing the computing devices **110** to perform particular operations at particular time or in a particular order.

[0075] As described above, data that results from performance of an extended memory operation can be stored in the original address in the computing device **110** in which the data was stored prior to performance of the extended memory operation, however, in some embodiments, blocks of data that result from performance of the extended memory operation can be converted into logical records subsequent to performance of the extended memory operation. The logical records can comprise data records that are independent of their physical locations. For example, the logical records may be data records that point to an address (e.g., a location) in at least one of the computing devices **110** where physical data corresponding to performance of the extended memory operation is stored.

[0076] In some embodiments, the result of the extended memory operation can be stored in an

address of a computing device memory that is the same as the address in which the data is stored prior to performance of the extended memory operation. Embodiments are not so limited, however, and the result of the extended memory operation can be stored in an address of the computing device memory that is the same as the address in which the data is stored prior to performance of the extended memory operation. In some embodiments, the logical records can point to these address locations such that the result(s) of the extended memory operation can be accessed from the computing devices **110** and transferred to circuitry external to the computing devices **110** (e.g., to a host).

[0077] In some embodiments, the controller **115** can receive and/or send blocks of data directly to and from the memory controller **112**. This can allow the controller **115** to transfer blocks of data that are not processed (e.g., blocks of data that are not used in performance of extended memory operations) by the computing devices **110** to and from the memory controller **112**.

[0078] For example, if the controller **115** receives unprocessed blocks of data from a host **102** coupled to the storage controller **104** that are to be stored by memory device(s) **116** coupled to the storage controller **104**, the controller **115** can cause the unprocessed blocks of data to be transferred to the memory controller **112**, which can, in turn, cause the unprocessed blocks of data to be transferred to memory device(s) coupled to the storage controller **104**.

[0079] Similarly, if the host requests an unprocessed (e.g., a full) block of data (e.g., a block of data that is not processed by the computing devices **110**), the memory controller **112** can cause unprocessed blocks of data to be transferred to the controller **115**, which can subsequently transfer the unprocessed blocks of data to the host.

[0080] FIG. 2 is a functional block diagram in the form of a computing system **200** including an apparatus **204** including a first plurality of communication subsystems **208**, a second plurality of communication subsystems **206**, and a plurality of memory devices **216** in accordance with a number of embodiments of the present disclosure. As used herein, an “apparatus” can refer to, but is not limited to, any of a variety of structures or combinations of structures, such as a circuit or circuitry, a die or dice, a module or modules, a device or devices, or a system or systems, for example. In the embodiment illustrated in FIG. 2, memory devices **216-1 . . . 216-N** can include one or more memory modules (e.g., double data rate (DDR) memory, three-dimensional (3D) cross-point memory, NAND memory, single in-line memory modules, dual in-line memory modules, etc.). The memory devices **216-1, . . . , 216-N** can include volatile memory and/or non-volatile memory. In a number of embodiments, memory devices **216-1, . . . , 216-N** can include a multi-chip device. A multi-chip device can include a number of different memory types and/or memory modules. For example, a memory system can include non-volatile or volatile memory on any type of a module.

[0081] As illustrated in FIG. 2, and in contrast to FIG. 1, a plurality of computing devices **210-1, 210-2, 210-3, 210-4, 210-5** (hereinafter referred to collectively as plurality of computing devices **210**) can be coupled to a first **208-1** of the first plurality of communication subsystem **208**, which is coupled to the plurality of hardware accelerators **114** through a second **206-2** of the second plurality of communication subsystems **206**. In one embodiment, the first plurality of communication subsystems **208** can be a plurality of multiplexers and the second plurality of communication subsystems **206** can be a plurality of AXI interconnects. Further, the first communication subsystem **208-1** is coupled directly to a buffer **219** which is coupled to the second communication subsystem **206-2**. The second **206-2** of the second plurality of communication subsystems **206** is coupled directly to an additional buffer **217**. The additional buffer **217** is coupled to a second **208-2** of the first plurality of communication subsystems **208**. The second **208-2** of the first plurality of communication subsystems **208** can be coupled to each of the plurality of hardware accelerators **214-1, 214-2, 214-3**. The hardware accelerators **214** can be on a same field programmable gate array (FPGA) as the computing devices **210**, first plurality of communication subsystems **208**, second plurality of communication subsystems **206**, etc. The hardware

accelerators **214** can be used for performing a number of posit operations, and/or for communication with an internal SRAM on the FPGA.

[0082] The first plurality of communication subsystems **208** can include circuitry and/or logic configured to allocate and de-allocate resources to the computing devices **210** during performance of operations described herein. For example, the circuitry and/or logic can allocate and/or de-allocate resources to the computing devices **210** during performance of extended memory operations described herein. While the examples described above include a particular number of multiplexers within a particular arrangement, examples are not so limited. For example, a multiplexer can be positioned between the buffer **219** and the second communication subsystem **206-2**, between the second communication subsystem **206-2** and the buffer **208-3**, etc. A third **208-3** of the first plurality of communication subsystems **208** can be coupled to a third of the second plurality of communication subsystems **206-3**. The third communication subsystem **206-3** can be coupled to a host interface **220**. In one example, the third communication subsystem **206-3** can be coupled to the host interface **220** via a number of AXI buses.

[0083] As is illustrated, a first (**206-1**) of the second plurality of communication subsystems **206** can be coupled to the controller (e.g., memory controller) **212**). The controller **212** can be coupled to a number of memory devices **216-1**, . . . , **216-N** via a number of channels **207-1**, . . . , **207-N**.

[0084] FIG. **3** is a functional block diagram in the form of a computing system **300** including an apparatus **304** including a plurality of communication subsystems **306**, **308** and a plurality of memory devices **316** in accordance with a number of embodiments of the present disclosure. As used herein, an “apparatus” can refer to, but is not limited to, any of a variety of structures or combinations of structures, such as a circuit or circuitry, a die or dice, a module or modules, a device or devices, or a system or systems, for example. In the embodiment illustrated in FIG. **3**, memory devices **316-1** . . . **316-N** can include one or more memory modules (e.g., double data rate (DDR) memory, three-dimensional (3D) cross-point memory, NAND memory, single in-line memory modules, dual in-line memory modules, etc.). The memory devices **316-1**, . . . , **316-N** can include volatile memory and/or non-volatile memory. In a number of embodiments, memory devices **316-1**, . . . , **316-N** can include a multi-chip device. A multi-chip device can include a number of different memory types and/or memory modules. For example, a memory system can include non-volatile or volatile memory on any type of a module.

[0085] As illustrated in FIG. **3**, the apparatus **304** can include a computing device (e.g., computing core). In some embodiments, the apparatus **304** can be an FPGA. As illustrated in FIG. **3**, the plurality of computing devices **310** can include ports **311** that can each be coupled to the plurality of communication subsystems **306** (as an example, without being coupled via an additional set of communication subsystems, such as communication subsystems **108** and **208**, (which may be multiplexers) illustrated in FIGS. **1** and **2**, respectively. The computing device **310** can be coupled to the plurality of communication subsystems **306** via corresponding port connections including a memory port (“MemPort”) **311-1**, system port (“SystemPort”) **311-2**, peripheral port (“PeriphPort”) **311-3**, and front port (“FrontPort”) **311-4**).

[0086] The memory port **311-1** can be directly coupled to a communication subsystem **306-1** specifically designated to receive data from a memory port and transfer the data to a memory controller **312**. The system port **311-2** can be directly coupled to a communication subsystem **308** that is further coupled to a plurality of buffers **319-1**, **319-2**, **319-3** (hereinafter referred to collectively as buffers **319**). Each of the plurality of buffers **319** can be coupled to a respective one of a plurality of communication subsystems **306-2**, **306-3**, **306-4**. The plurality of communication subsystems **306-2**, **306-3**, **306-4** can be coupled to an additional plurality of buffers **317-1**, **317-2**, **317-3**. The plurality of buffers **317** are each coupled to a respective one of a plurality of hardware accelerators **314-1**, **314-2**, **314-3**. The plurality of hardware accelerators **314** are coupled to logic **313**. The plurality of communication subsystems **306-2**, **306-3**, **306-4** are each specifically designated to receive data from the system port **311-2** and transfer the data to a respective

accelerator (e.g., an on-chip accelerator) **314**, which can then transfer data to additional logic circuitry **313**.

[0087] The peripheral port **311-3** can be directly coupled to a communication subsystem **306-5** specifically designated to receive data from the peripheral port **311-3** and transfer the data to a serial port **318**. The front port **311-4** can be directly coupled to a communication subsystem **306-6** specifically designated to receive data from the front port **311-4** and transfer the data to a host interface **320**, and subsequently to a host **302** via channels **303** and/or **305**. In this embodiment, the hardware accelerators **314** may be coupled to the computing device **310** via a multiplexer. In contrast, a multiplexer may not be used to couple the controller **312**, the serial port **318**, and/or the host interface **320** to the computing device **310**, but rather the ports and the communication subsystem are directly connected for data transfer.

[0088] In some embodiments, the communication subsystems **306** can facilitate visibility between respective address spaces of the computing device **310**. For example, the computing device **310** can, responsive to receipt of data and/or a file, store the data in a memory resource of the computing device **310**. The computing device can associate an address (e.g., a physical address) corresponding to a location in the memory resource of the computing device **310** in which the data is stored. In addition, the computing device **310** can parse (e.g., break) the address associated with the data into logical blocks.

[0089] In some embodiments, the zeroth logical block associated with the data can be transferred to a processing device (e.g., a reduced instruction set computing (RISC) device). A particular computing device (e.g., computing device **110**, **210**, **310**) can be configured to recognize that a particular set of logical addresses are accessible to that computing device (e.g., **210-2**), while other computing devices (e.g., computing device **210-3**, **210-4**, respectively, etc.) can be configured to recognize that different sets of logical addresses are accessible to those computing devices **110**, **210**, **310**. Stated alternatively, a first computing device (e.g., the computing device **210-2**) can have access to a first set of logical addresses associated with that computing device (**210-2**), and a second computing device (e.g., the computing device **210-3**) can have access to a second set of logical address associated therewith, etc.

[0090] If data corresponding to the second set of logical addresses (e.g., the logical addresses accessible by the second computing device **210-3**) is requested at the first computing device (e.g., the computing device **210-2**), the communication subsystems **306** can facilitate communication between the first computing device (e.g., the computing device **210-2**) and the second computing device (e.g., the computing device **210-3**) to allow the first computing device (e.g., the computing device **210-2**) to access the data corresponding to the second set of logical addresses (e.g., the set of logical addresses accessible by the second computing device **210-3**). That is, the communication subsystem **308** can facilitate communication between the computing device **310** (e.g., **210-1**) and additional computing devices (e.g., computing devices **210-2**, **210-3**, **210-4**) to allow address spaces of the computing devices to be visible to one another.

[0091] In some embodiments, communication between the computing devices **110**, **210**, **310** to facilitate address visibility can include receiving, by an event queue of the first computing device (e.g., the computing device **210-1**), a message requesting access to the data corresponding to the second set of logical addresses, loading the requested data into a memory resource of the first computing device, and transferring the requested data to a message buffer. Once the data has been buffered by the message buffer, the data can be transferred to the second computing device (e.g., the computing device **210-2**) via the communication subsystem **310**.

[0092] For example, during performance of an extended memory operation, the controller **115**, **215**, **315** and/or a first computing device (e.g., the computing device **210-1**) can determine that the address specified by a host command (e.g., a command to initiate performance of an extended memory operation generated by a host such as the host **102** illustrated in FIG. 1) corresponds to a location in a memory resource of a second computing device (e.g., the computing device **210-2**)

among the plurality of computing devices (**110, 210**). In this case, a computing device command can be generated and sent from the controller **115, 215, 315** and/or the first computing device (**210-1**) to the second computing device (**210-2**) to initiate performance of the extended memory operation using an operand stored in the memory resource of the second computing device (**210-2**) at the address specified by the computing device command.

[0093] In response to receipt of the computing device command, the second computing device (**210-2**) can perform the extended memory operation using the operand stored in the memory resource of the second computing device (**210-2**) at the address specified by the computing device command. This can reduce command traffic from between the host and the storage controller and/or the computing devices (**210, 310**), because the host need not generate additional commands to cause performance of the extended memory operation, which can increase overall performance of a computing system by, for example reducing a time associated with transfer of commands to and from the host.

[0094] In some embodiments, the controller **115, 215, 315** can determine that performing the extended memory operation can include performing multiple sub-operations. For example, an extended memory operation may be parsed or broken into two or more sub-operations that can be performed as part of performing the overall extended memory operation. In this case, the controller **115, 215, 315** and/or the communication subsystems (**106, 108, 206, 208, 308**) can utilize the above described address visibility to facilitate performance of the sub-operations by various computing devices **110, 210, 310**. In response to completion of the sub-operation, the controller **115, 215, 315** can cause the results of the sub-operations to be coalesced into a single result that corresponds to a result of the extended memory operation.

[0095] In other embodiments, an application requesting data that is stored in the computing devices **110, 210, 310** can know (e.g., can be provided with information corresponding to) which computing devices **110, 210, 310** include the data requested. In this example, the application can request the data from the relevant computing device **110, 210, 310** and/or the address may be loaded into multiple computing devices **110, 210, 310** and accessed by the application requesting the data via the communication subsystems **108, 106, 208, 206, 308**.

[0096] The controller **115, 215, 315** can be discrete circuitry that is physically separate from the communication subsystems **108, 106, 208, 206, 308** and can each be provided as one or more integrated circuits that allows communication between the computing devices **110, 210, 310**, the memory controller **112, 212, 312** and/or the controller **115, 215, 315**. Non-limiting examples of communication subsystems **108, 106, 208, 206, 308** can include a XBAR or other communications subsystem that allows for interconnection and/or interoperability of the controller **115, 215, 315**, the computing devices **110, 210, 310**, and/or the memory controller **112, 212, 312**.

[0097] As described above, responsive to receipt of a command generated by the controller **115, 215, 315**, the communication subsystems **108, 106, 208, 206, 308**, and/or a host (e.g., the host **102** illustrated in FIG. 1), performance of extended memory operations using data stored in the computing devices **110, 210, 310** and/or from blocks of data streamed through the computing devices **110, 210, 310** can be realized.

[0098] FIG. 4 is a functional block diagram in the form of a computing system **400** including an apparatus **404** including a first plurality of communication subsystems **406**, a second communication subsystem **408**, and a plurality of memory devices **416** in accordance with a number of embodiments of the present disclosure.

[0099] As illustrated in FIG. 4, the apparatus **404** can include a computing device (e.g., computing core). In some embodiments, the apparatus **404** can be an FPGA. As illustrated in FIG. 4, and similarly in FIG. 3, the plurality of computing devices **410** can include ports **411** that can each be coupled to the plurality of communication subsystems **406** (as an example, without being coupled via an additional set of communication subsystems, such as communication subsystems **108** and **208**, (which may be multiplexers) illustrated in FIGS. 1 and 2, respectively. The computing device

410 can be coupled to the plurality of communication subsystems **406** via corresponding port connections including a memory port (“MemPort”) **411-1**, system port (“SystemPort”) **411-2**, peripheral port (“PeriphPort”) **411-3**, and front port (“FrontPort”) **411-4**).

[0100] The memory port **411-1** can be directly coupled to a communication subsystem **406-1** specifically designated to receive data from a memory port and transfer the data to a memory controller **412**. In contrast to FIG. 3, FIG. 4 illustrates the system port **411-2** being directly coupled to a buffer **419** that is coupled directly to the second communication subsystem **406-2**. The second communication subsystem **406-2** is likewise coupled to an additional buffer **417**. The additional buffer **417** is coupled to an additional communication subsystem **408**. As is shown similarly in FIG. 2, FIG. 4 illustrates a communication subsystem **408** that is coupled to each of a plurality of hardware accelerators **414-1**, **414-2**, **414-3**, respectively.

[0101] Further, the peripheral port **411-3** can be directly coupled to a communication subsystem **406-5** specifically designated to receive data from the peripheral port **411-3** and transfer the data to a serial port **418**. The front port **411-4** can be directly coupled to a communication subsystem **406-6** specifically designated to receive data from the front port **411-4** and transfer the data to a host interface **420**, and subsequently to a host **402** via channels **403** and/or **405**. In this embodiment, the hardware accelerators **414** may be coupled to the computing device **410** via a multiplexer. In contrast, a multiplexer may not be used to couple the controller **412**, the serial port **418**, and/or the host interface **420** to the computing device **310**, but rather the ports and the communication subsystem are directly connected for data transfer.

[0102] FIG. 5 is a functional block diagram in the form of a computing core **510** including a number of ports **511-1**, **511-2**, **511-3**, **511-4** in accordance with a number of embodiments of the present disclosure. The computing core **510** can include a memory management unit (MMU) **520**, a physical memory protection (PMP) unit **522**, and a cache **524**.

[0103] The MMU **520** refers to a computer hardware component used for memory and caching operations associated with a processor. The MMU **520** can be responsible for memory management and be integrated into the processor, or, in some examples, can be on a separate integrated circuit (IC) chip. The MMU **520** can be used for hardware memory management which can include overseeing and regulating the processor's use of random access memory (RAM) and cache memory. The MMU **520** can be used for operating system (OS) memory management, which can ensure the availability of adequate memory resources for the objects and data structures of each running program. The MMU **520** can be used for application memory management, which can allocate each individual program's required or used memory, and then recycle freed up memory space when the operation concludes or the space becomes available.

[0104] In one embodiment, the PMP unit **522** can be used to restrict access to memory and isolate processes from each other. The PMP unit **522** can be used to set memory access privileges (read, write, execute) for specified memory regions. The PMP unit **522** can support 8 regions with a minimum region size of 4 bytes. In some examples, the PMP unit **522** may only be programmed in M-mode. The PMP unit **522** may enforce permissions on U-mode accesses. However, locked regions may additionally enforce their permissions on M-mode. The cache **524** can be an SRAM cache, a 3D cross-point cache, etc. The cache **524** can include 8 KB, 16 KB, 32 KB, etc. and can include error correction coding (ECC).

[0105] The computing core **510** can also include a plurality of ports including a memory port **511-1**, a system port **511-2**, a peripheral port **511-3**, and a front port **511-4**. The memory port **511-1** can be directly coupled to a communication subsystem (as illustrated in FIG. 3) specifically designated to receive data from a memory port **511-1**. The system port **511-2** can be directly coupled to a communication subsystem specifically designated to receive data from the system port **511-2**. The data through the system port **511-2** can be transferred to an accelerator (e.g., an on-chip accelerator). The peripheral port **511-3** can be directly coupled to a communication subsystem specifically designated to receive data from the peripheral port **511-3** and this data can be

eventually transferred to a serial port. The front port **511-4** can be directly coupled to a communication subsystem specifically designated to receive data from the front port **511-4** and this data can be eventually transferred to a host interface, and subsequently to a host.

[0106] The computing core **510** can be a full-Linux capable, cache-coherent 64-bit RISC-V processor. In some examples, the memory port **511-1**, the system port **511-2**, and the peripheral port **511-3** can be outgoing ports and the front port **511-4** can be an incoming port. An example of computing core **510** can include a U54-MC computing core. The computing core **510** can include an instruction memory system, an instruction fetch unit, an execution pipeline unit, a data memory system, and support for global, software, and timer interrupts. The instruction memory system can include a 16 Kibibyte (KiB) 2-way set-associative instruction cache. The access latency of all blocks in the instruction memory system can be one clock cycle. The instruction cache may not be kept coherent with the rest of the platform memory system. Writes to the instruction memory may be synchronized with the instruction fetch stream by executing a FENCE.I instructions. The instruction cache can have a line size of 64 bytes, and a cache line fill can trigger a burst access outside the computing core **510**.

[0107] The instruction fetch unit can include branch prediction hardware to improve performance of the processor core. The branch predictor can include a 28-entry branch target buffer (BTB), which can predict a target of taken branches, a 512-entry branch history table (BHT), which can predict the direction of conditional branches, and a 6-entry return-address stack (RAS) which can predict a target of procedure returns. The branch predictor may have one-cycle latency, so that correctly predicted control-flow instructions result in no penalty. An incorrect prediction of control-flow instructions may incur three-cycle penalty.

[0108] The execution pipeline unit can be a single-issue, in-order pipeline. The pipeline can include five stages: instruction fetch, instruction decode and register fetch, execute, data memory access, and register writeback. The pipeline can have a peak execution rate of one instruction per clock cycle, and may be fully bypassed so that most instructions have a one-cycle result latency. The pipeline may interlock on read-after-write and write-after-write hazards, so instructions may be scheduled to avoid stalls.

[0109] The data memory system can include a DTIM interface, which can support up to 8 KiB. The access latency from a core to its own DTIM may be two clock cycles for full words and three clock cycles for smaller quantities. Memory requests from one core to any other core's DTIM may not be as performant as memory requests from a core to its own DTIM. Misaligned accesses are not supported in hardware and may result in a trap to allow software emulation.

[0110] In some embodiments, the computing core **510** can include a floating-point unit (FPU) which can provide full hardware support for the IEEE 754-2008 floating-point standard for 32-bit single-precision and 64-bit double-precision arithmetic. The FPU can include a fully pipelined fused-multiply-add unit and an iterative divide and square-root unit, magnitude comparators, and float-to-integer conversion units, with full hardware support for subnormals and IEEE default values.

[0111] FIG. **6** is a flow diagram representing an example method **628** corresponding to an extended memory interface in accordance with a number of embodiments of the present disclosure. At block **630**, the method **628** can include receiving, at a processing unit that is coupled between a host device and a non-volatile memory device, signaling indicative of a plurality of operations to be performed on data written to or read from the non-volatile memory device. The quantity of operations can include extended memory operations as described above.

[0112] At block **632**, the method **628** can include performing, at the processing unit, at least one operation of the plurality of operations in response to the signaling. A computing device (such as computing device **110**, **210**, **310**, **410** in FIGS. **1-4**, respectively) can include the processing unit that performs the at least one operation. The operation can be performed using a block of data in response to receipt of the block of data to reduce a size of data from a first size to a second size by

the at least one of the plurality of computing devices. The performance of the operation can be caused by a controller. The controller can be analogous to the controller **115**, **215**, **315** illustrated in FIGS. **1-3**, herein. In some embodiments, performing the operation can include performing an extended memory operation, as described herein. The operation can further include performing, by the particular computing device, the operation in the absence of receipt of a host command from a host coupleable to the controller. In response to completion of performance of the operation, the method **628** can include sending a notification to a host coupleable to the controller.

[0113] At block **634**, the method **628** can include accessing, via a controller at the processing unit or non-volatile memory device, a portion of a memory array in the non-volatile memory device. The non-volatile memory device can be accessed by a memory controller and the memory controller can send the accessed data to a computing device, a hardware accelerator, etc. in order to perform one of the quantity of operations. The method **628** can further include causing, using an additional controller (e.g., memory controller), the blocks of data to be transferred from the memory device to a plurality of communication subsystems. The method **628** can further include allocating, via the pluralities of communication subsystems, resources corresponding to respective computing devices among the plurality of computing devices to perform the operation on the block of data.

[0114] At block **636**, the method **628** can include transmitting, to a hardware accelerator, additional signaling indicative of a command to perform one or more additional operations of the plurality of operations on the data written to or read from the non-volatile memory device. For example, signaling indicative of a first operation can be sent to a first hardware accelerator, signaling indicative of a second operation can be sent to a second hardware accelerator, etc.

[0115] In some embodiments, the command to initiate performance of the operation can include an address corresponding to a location in the memory array of the particular computing device and the method **628** can include storing a result of the operation in the address corresponding to the location in the particular computing device. For example, the method **628** can include storing a result of the operation in the address corresponding to the memory location in the particular computing device in which the operand corresponding to performance of the operation was stored prior to performance of the extended memory operation. That is, in some embodiments, a result of the operation can be stored in the same address location of the computing device in which the data that was used as an operand for the operation was stored prior to performance of the operation.

[0116] In some embodiments, the method **628** can include determining, by the controller, that the operand corresponding to performance of the operation is not stored by the particular computing device. In response to such a determination, the method **628** can further include determining, by the controller, that the operand corresponding to performance of the operation is stored in a memory device coupled to the plurality of computing devices. The method **628** can further include retrieving the operand corresponding to performance of the operation from the memory device, causing the operand corresponding to performance of the operation to be stored in at least one computing device among the plurality of computing device, and/or causing performance of the operation using the at least one computing device. The memory device can be analogous to the memory devices **116** illustrated in FIG. **1**.

[0117] The method **628** can, in some embodiments, further include determining that at least one sub-operation is to be performed as part of the operation, sending a command to a computing device different than the particular computing device to cause performance of the sub-operation, and/or performing, using the computing device different than the particular computing device, the sub-operation as part of performance of the operation. For example, in some embodiments, a determination that the operation is to be broken into multiple sub-operations can be made and the controller can cause different computing devices to perform different sub-operations as part of performing the operation. In some embodiments, the controller can, in concert with the first and the second pluralities of communications subsystem, such **108**, **106**, **208**, **206**, **308**, **306**, and **408**, **406**

illustrated in FIGS. 1-4, herein, assign sub-operations to two or more of the computing devices as part of performance of the operation and/or to two or more of the hardware accelerators.

[0118] Although specific embodiments have been illustrated and described herein, those of ordinary skill in the art will appreciate that an arrangement calculated to achieve the same results can be substituted for the specific embodiments shown. This disclosure is intended to cover adaptations or variations of one or more embodiments of the present disclosure. It is to be understood that the above description has been made in an illustrative fashion, and not a restrictive one. Combination of the above embodiments, and other embodiments not specifically described herein will be apparent to those of skill in the art upon reviewing the above description. The scope of the one or more embodiments of the present disclosure includes other applications in which the above structures and processes are used. Therefore, the scope of one or more embodiments of the present disclosure should be determined with reference to the appended claims, along with the full range of equivalents to which such claims are entitled.

[0119] In the foregoing Detailed Description, some features are grouped together in a single embodiment for the purpose of streamlining the disclosure. This method of disclosure is not to be interpreted as reflecting an intention that the disclosed embodiments of the present disclosure have to use more features than are expressly recited in each claim. Rather, as the following claims reflect, inventive subject matter lies in less than all features of a single disclosed embodiment. Thus, the following claims are hereby incorporated into the Detailed Description, with each claim standing on its own as a separate embodiment.

Claims

1. An apparatus, comprising: a plurality of memory devices; and a computing block coupled to the plurality of memory devices and comprising: a number of computing devices; and a number of hardware accelerators; and wherein the computing block is configured to: receive, from a host, a command to perform an operation using data read from a memory device of the plurality of memory devices; transmit sub-operations of the operation to respective different hardware accelerators of the number of hardware accelerators; execute, at the respective different hardware accelerators, the sub-operations; and generate a result of the operation by combining results of the sub-operations.
2. The apparatus of claim 1, wherein the number of computing devices comprise RISC processors.
3. The apparatus of claim 1, wherein the command is received at a particular computing device of the number of computing devices, and the result of the operation is generated at the particular computing device.
4. The apparatus of claim 1, wherein the computing block is a field programmable gate array (FPGA).
5. The apparatus of claim 4, wherein the computing block is a controller.
6. The apparatus of claim 1, wherein the plurality of memory devices are volatile memory devices.
7. The apparatus of claim 1, wherein the plurality of memory devices comprise volatile memory devices and nonvolatile memory devices.
8. The apparatus of claim 1, wherein at least one sub-operation is provided to one of the number hardware accelerators from a different one of the number of hardware accelerators.
9. The apparatus of claim 1, wherein the sub-operations of the operation are transmitted to the respective different hardware accelerators from at least two computing devices of the number of computing devices.
10. The apparatus of claim 1, wherein the results of the sub-operations are combined by one of the hardware accelerators to which the sub-operations were transmitted.
11. An apparatus, comprising: a memory device; and a computing block coupled to the memory device and comprising a computing device and a plurality of hardware accelerators, wherein the

computing device is configured to: receive a command to perform an operation using data stored in the memory device; divide the operation into a number of sub-operations; send a first sub-operation to a first hardware accelerator for execution; send a second sub-operation to a second hardware accelerator for execution; and combine results of the first and second sub-operations into a result of the operation.

12. The apparatus of claim 11, wherein the plurality of hardware accelerators are coupled to respective dedicated buffers.

13. The apparatus of claim 11, wherein the plurality of hardware accelerators are coupled to logic circuitry external to the computing block.

14. The apparatus of claim 11, wherein the operation involves reducing a size of a block of data from a first size to a second size.

15. The apparatus of claim 11, wherein the operation is an extended memory operation that involves the data being reordered, removed, or discarded.

16. The apparatus of claim 11, further comprising a host external to the computing block and coupled thereto via a host interface.

17. The apparatus of claim 16, wherein the host interface is a double data rate (DDR) interface.

18. An apparatus, comprising: a first hardware accelerator and a second hardware accelerator; and a processing unit configured to: receive a command to perform an operation on data stored in a memory device to which the processing device is coupled; and send the command to the first hardware accelerator for execution, wherein the first hardware accelerator is configured to: execute a first portion of the operation to generate a first result; send a second portion of the operation to the second hardware accelerator; and upon receipt, from the second hardware accelerator, of a second result of the second portion of the operation, generate a result of the operation by combining the first result and the second result.

19. The apparatus of claim 18, wherein the processing unit and the first and second hardware accelerators are located on a controller located external to a host from which the command is received.

20. The apparatus of claim 18, wherein the processing unit is one of a plurality of processing units having ports coupled to respective communication subsystems.
