



US 20250267129A1

(19) **United States**

(12) **Patent Application Publication**  
**Eshghi**

(10) **Pub. No.: US 2025/0267129 A1**

(43) **Pub. Date: Aug. 21, 2025**

(54) **PREDICTING USER-FILE INTERACTIONS**

**Publication Classification**

(71) Applicant: **Box, Inc.**, Redwood City, CA (US)

(72) Inventor: **Kave Eshghi**, Los Altos, CA (US)

(73) Assignee: **Box, Inc.**, Redwood City, CA (US)

(21) Appl. No.: **19/031,303**

(22) Filed: **Jan. 17, 2025**

(51) **Int. Cl.**

**H04L 9/40** (2022.01)

**G06F 16/901** (2019.01)

**G06N 7/01** (2023.01)

(52) **U.S. Cl.**

CPC ..... **H04L 63/029** (2013.01); **G06F 16/9024**

(2019.01); **G06N 7/01** (2023.01); **H04L**

**63/1425** (2013.01)

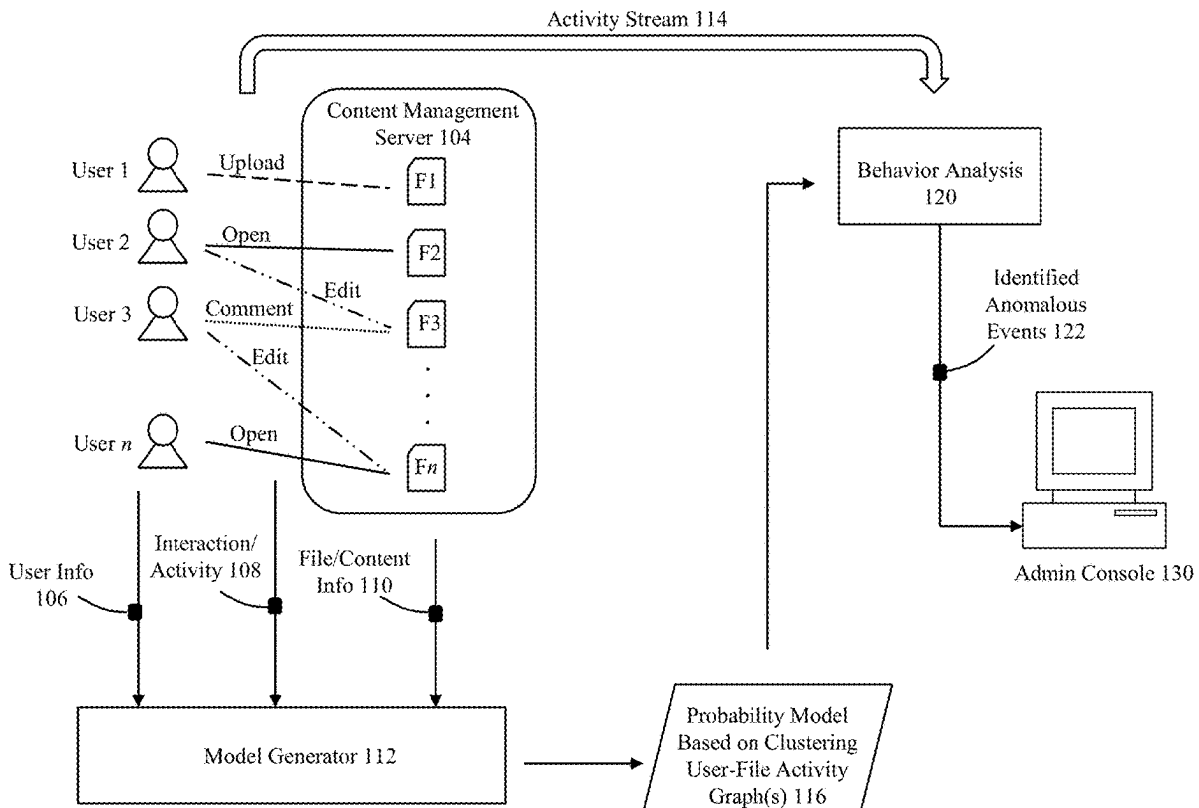
**Related U.S. Application Data**

(63) Continuation of application No. 18/492,727, filed on Oct. 23, 2023, now Pat. No. 12,309,124, which is a continuation of application No. 17/322,360, filed on May 17, 2021, now Pat. No. 11,799,835, which is a continuation of application No. 16/115,069, filed on Aug. 28, 2018, now Pat. No. 11,012,421.

(57)

**ABSTRACT**

Disclosed is an improved systems, methods, and computer program products that use a cluster-based probability model to perform anomaly detection, where the clusters are based upon entities and interactions that exist in content management platforms.



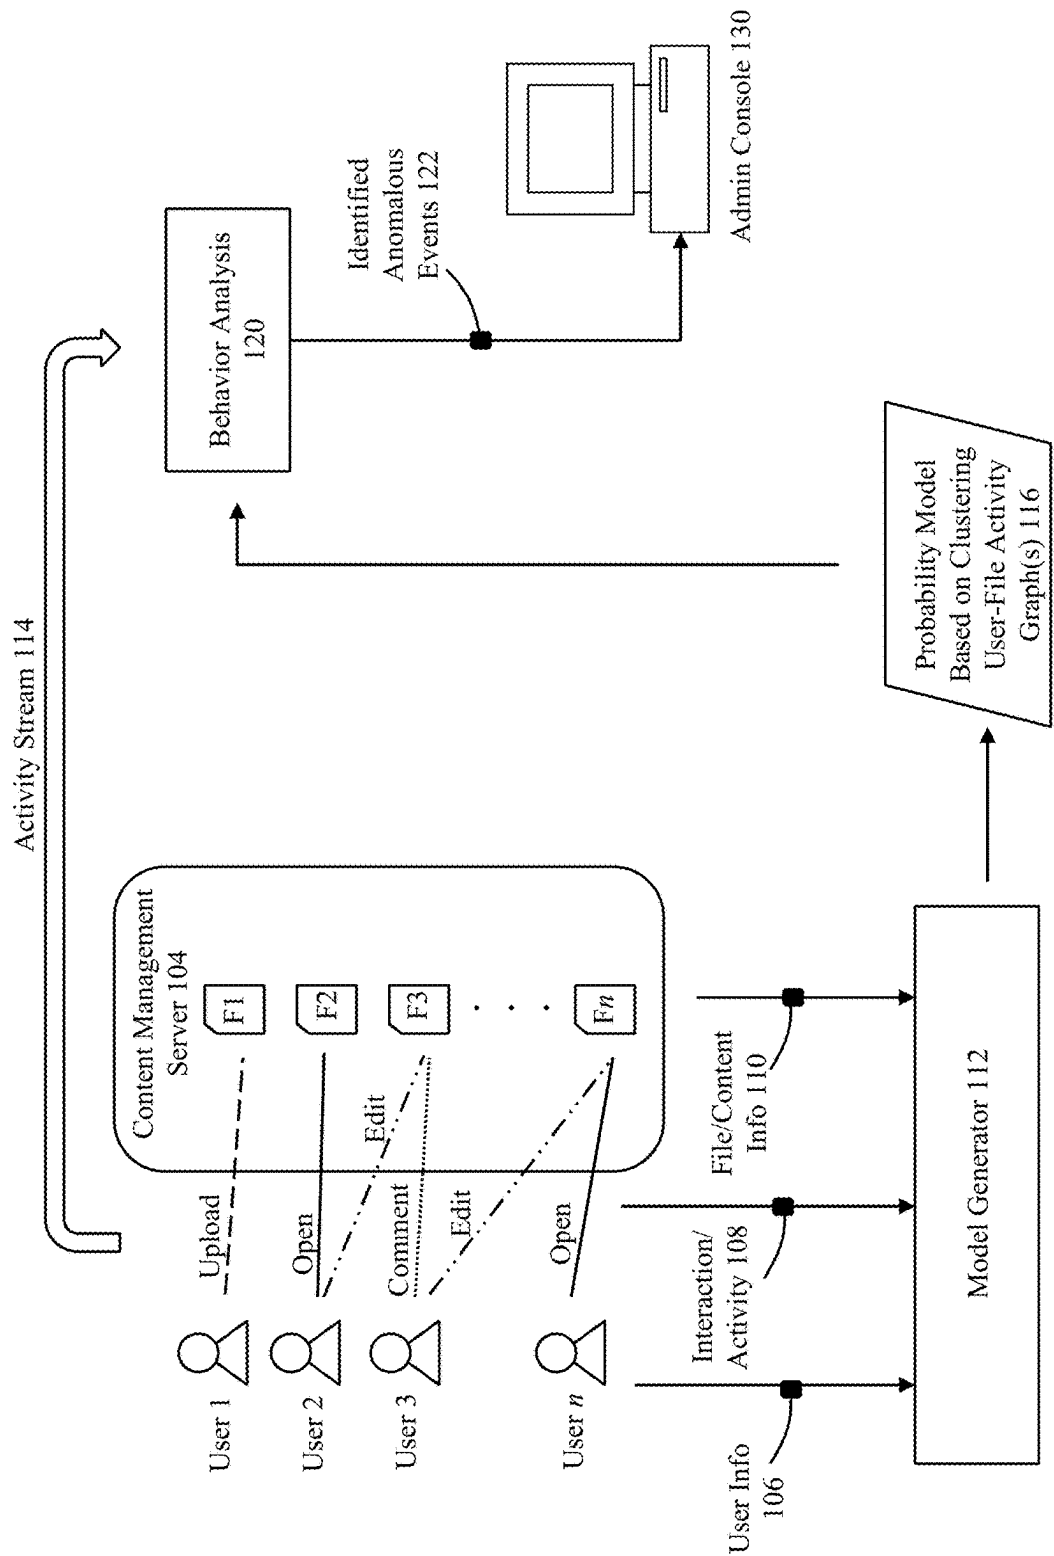


Figure 1




Figure 2

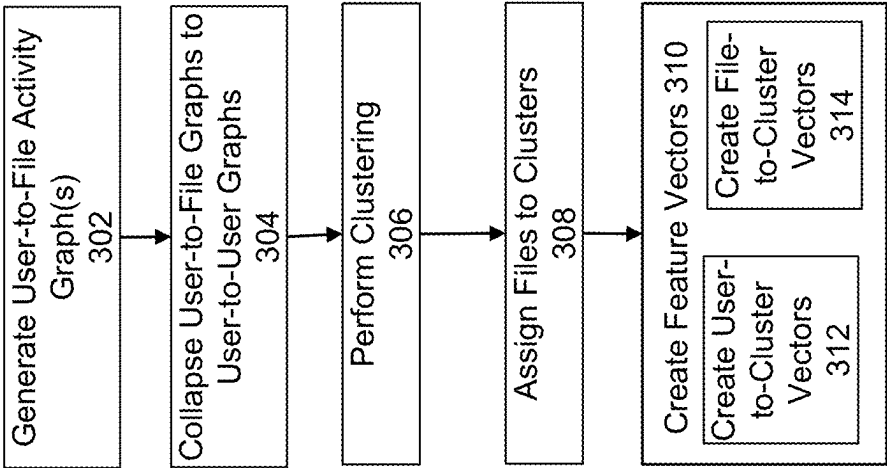


Figure 3




Figure 4




Figure 5




Figure 6

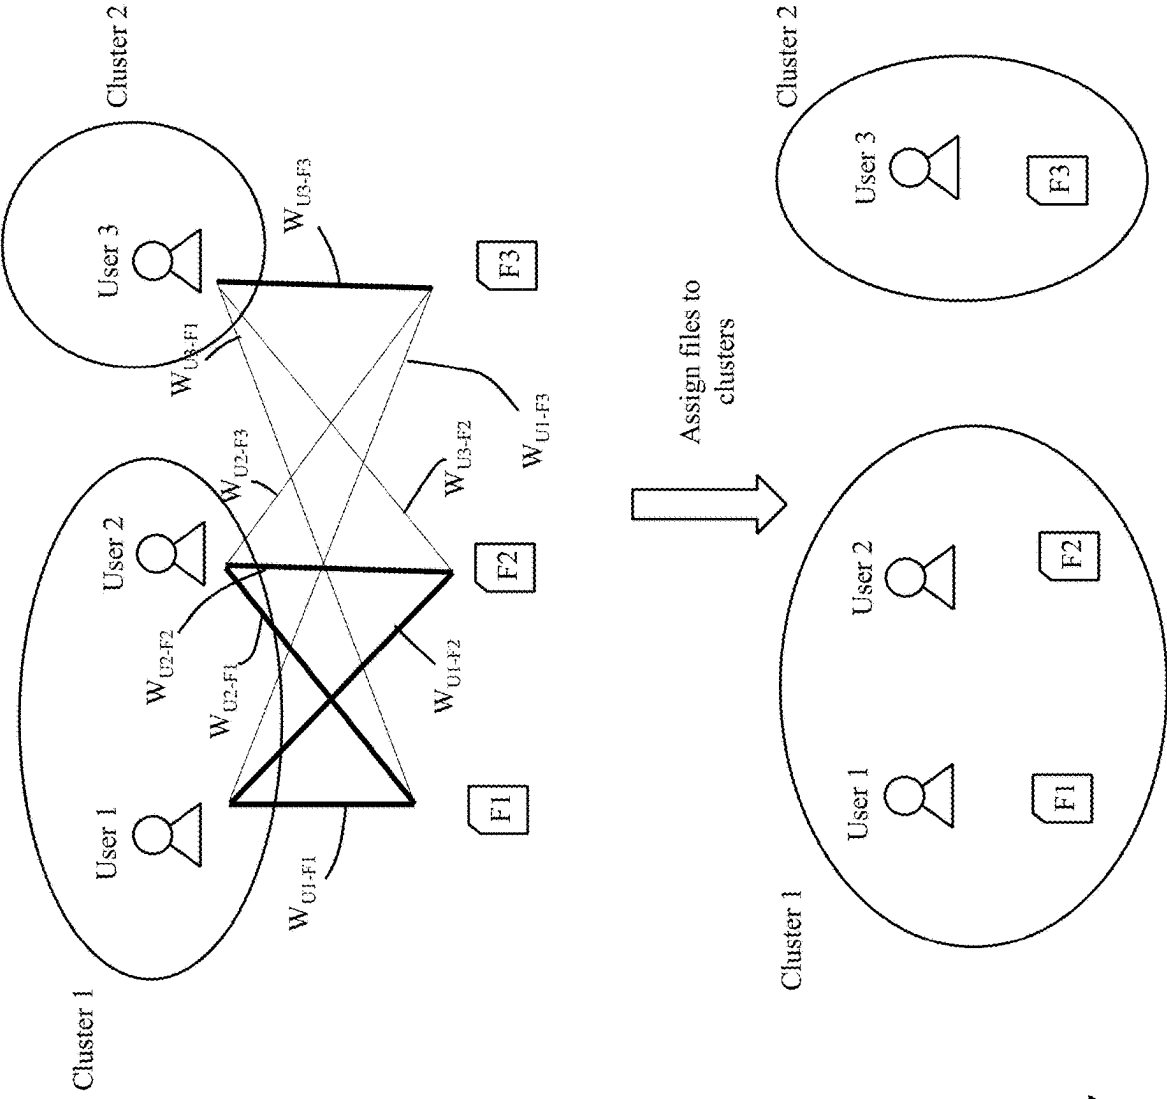


Figure 7






Figure 8




Figure 9




Figure 10

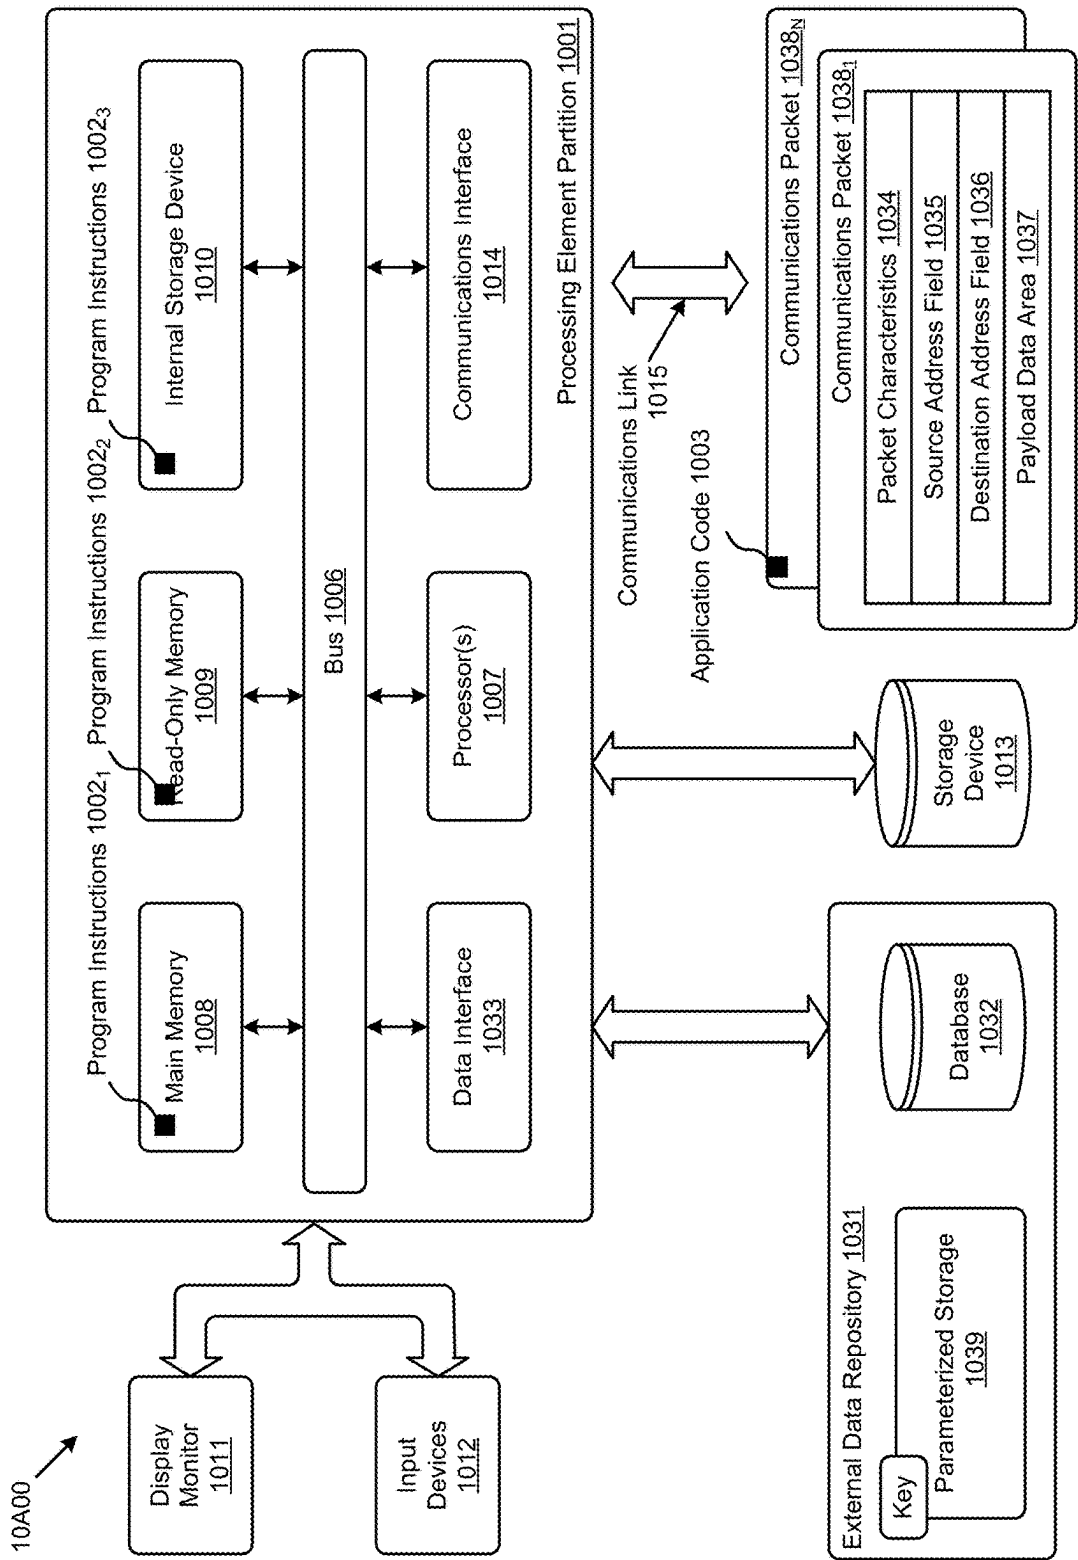



FIG. 11A



**PREDICTING USER-FILE INTERACTIONS****CROSS-REFERENCE TO RELATED U.S.  
PATENT APPLICATION(S)**

[0001] The present application is a continuation of U.S. patent application Ser. No. 18/492,727, entitled “PREDICTING USER-FILE INTERACTIONS”, filed on Oct. 23, 2023, which is a continuation of U.S. patent application Ser. No. 17/322,360, filed on May 17, 2021, issued on Oct. 24, 2023, as U.S. Pat. No. 11,799,835, which is a continuation of U.S. patent application Ser. No. 16/115,069 entitled “PREDICTING USER-FILE INTERACTIONS”, filed on Aug. 28, 2018, issued on May 18, 2021, as U.S. Pat. No. 11,012,421, which are hereby incorporated by reference in their entirety.

**BACKGROUND**

[0002] The proliferation of cloud-based services and platforms continues to increase. Specifically, cloud-based content management services and platforms have impacted the way personal and corporate information is stored, and has also impacted the way personal and corporate information is shared and managed. Individuals can avail themselves of such storage and management services to store and share large volumes of content objects such as files and documents.

[0003] Enterprises (e.g., companies, firms, etc.) might also want to use cloud-based content management platforms to secure, store and manage proprietary content objects, while enhancing the experience and productivity of their workforce (e.g., active employees) and/or their customers (e.g., clients, vendors, contractors, partners, etc.). For example, the enterprise might want to provide their active employees access to the flexible provisioning (e.g., access, roles, sharing, etc.) and collaboration components (e.g., tasks, comments, workflow, etc.) delivered by a content management platform provider.

[0004] While a key benefit of using such cloud-based systems is the ability to easily share content among users, issues may arise when the cloud-based system is employed to store and share content that is considered sensitive or proprietary (e.g., strategic documents, product specifications, financial statements, etc.), and where dangerous user behaviors may occur with respect to access of the files in the cloud-based system. It is highly desirable to be able to detect if files are being accessed in ways that are indicative of external infiltration, disgruntled employees, stolen credentials, or the like.

[0005] Various legacy techniques have been implemented to provide security to documents. For example, many systems implement security at the firewall level that seek to prevent improper access by third parties. However, such systems, while useful to prevent certain types of known attacks, may become useless if the system is successfully penetrated by a third party that are now accessing the enterprise's confidential files. Other legacy techniques may employ rules, filters, or learning algorithms to check for harmful behavior. However, conventional solutions are not specifically designed or optimized to address the unique types of accesses and interactions that occur in cloud-based content management platforms.

[0006] The problem to be solved is therefore rooted in technological limitations of the legacy approaches.

Improved techniques, in particular improved application of technology, are needed to address the problem of implementing document and file security in content management platforms.

[0007] What is needed is a technique or techniques to improve the application and efficacy of various technologies as compared with the application and efficacy of legacy approaches.

**SUMMARY**

[0008] The present disclosure provides improved systems, methods, and computer program products suited to address the aforementioned issues with legacy approaches. More specifically, the present disclosure provides a detailed description of techniques used in systems, methods, and in computer program products for using a cluster-based probability model to perform anomaly detection, where the clusters are based upon entities and interactions that exist in content management platforms.

[0009] Further details of aspects, objectives, and advantages of the disclosure are described below and in the detailed description, drawings, and claims. Both the foregoing general description of the background and the following detailed description are exemplary and explanatory, and are not intended to be limiting as to the scope of the claims.

**BRIEF DESCRIPTION OF THE DRAWINGS**

[0010] The drawings described below are for illustration purposes only. The drawings are not intended to limit the scope of the present disclosure.

[0011] FIG. 1 shows a system to perform anomaly detection according to some embodiments of the invention.

[0012] FIG. 2 shows a flowchart of an approach to perform anomaly detection according to some embodiments of the invention.

[0013] FIG. 3 shows a flowchart of an approach to generate a probability model according to some embodiments of the invention.

[0014] FIG. 4 illustrates formation of user-to-file graphs according to some embodiments of the invention.

[0015] FIG. 5 illustrates formation of user-to-user graphs according to some embodiments of the invention.

[0016] FIG. 6 illustrates clustering according to some embodiments of the invention.

[0017] FIG. 7 illustrates association of files to clusters according to some embodiments of the invention.

[0018] FIG. 8 illustrates example feature vectors according to some embodiments of the invention.

[0019] FIG. 9 shows a flowchart of an approach to generate scoring to perform behavioral analysis according to some embodiments of the invention.

[0020] FIG. 10 illustrates scoring according to some embodiments.

[0021] FIG. 11A and FIG. 11B depict exemplary architectures of components suitable for implementing embodiments of the present disclosure, and/or for use in the herein-described environments.

**DETAILED DESCRIPTION**

[0022] Some embodiments of the present disclosure provide improved systems, methods, and computer program products suited to address the aforementioned issues with legacy approaches for anomaly detection. More specifically,

the present disclosure provides a detailed description of techniques used in systems, methods, and in computer program products for using a cluster-based probability model to perform anomaly detection, where the clusters are based upon entities and interactions that exist in content management platforms.

**[0023]** FIG. 1 illustrates an example diagram of a system having a content management server **104** of a cloud service (e.g., collaboration or file sharing platform) and/or cloud storage accounts with capabilities for enabling a desktop application such as described herein for accessing a cloud-based collaboration platform. The system depicted in FIG. 1 may be broadly referred to as a cloud-based collaboration platform.

**[0024]** One or more users may use one or more client devices to access the content management server **104**. The client devices can be any system and/or device, and/or any combination of devices/systems that is able to establish a connection, including wired, wireless, cellular connections with another device, a server and/or other systems such as server **104**. The client devices will typically include a display and/or other output functionalities to present information and data exchanged between among the devices and/or the server **104**. For example, the client devices can include mobile, hand held or portable devices or non-portable devices and can be any of a server desktop, a desktop computer, a computer cluster, or portable devices including, a notebook, a laptop computer, a handheld computer, a palmtop computer, a mobile phone, a cell phone, a smart phone, a PDA, a Blackberry device, a handheld tablet (e.g. an iPad, a Galaxy, Xoom Tablet, etc.), a tablet PC, a thin-client, a hand held console, a hand held gaming device or console, an iPhone, and/or any other portable, mobile, hand held devices, etc. running on any platform or any operating system (e.g., Mac-based OS (OS X, IOS, etc.), Windows-based OS (Windows Mobile, Windows, etc.), Android, Blackberry OS, Embedded Linux platforms. In one embodiment, the client devices and the server **104** are coupled via a network. In some embodiments, the devices and server **104** may be directly connected to one another. The input mechanism on the client devices can include touch screen keypad (including single touch, multi-touch, gesture sensing in 2D or 3D, etc.), a physical keypad, a mouse, a pointer, a track pad, motion detector (e.g., including 1-axis, 2-axis, 3-axis accelerometer, etc.), a light sensor, capacitance sensor, resistance sensor, temperature sensor, proximity sensor, a piezoelectric device, device orientation detector (e.g., electronic compass, tilt sensor, rotation sensor, gyroscope, accelerometer), or a combination of the above.

**[0025]** Signals received or detected indicating user activity at the client devices through one or more of the above input mechanism, or others, can be used in the disclosed technology by any user or collaborator (hereinafter referred to collectively as “user”) for accessing a cloud-based collaboration platform or online collaboration platform (e.g., hosted by the host server **104**).

**[0026]** The collaboration platform or environment hosts workspaces with content items that one or more users can access (e.g., view, edit, update, revise, comment, download, preview, tag, or otherwise manipulate, etc.). To host the content items, the collaboration platform may store the uploaded content items in a data repository. The data repository may represent the collective data storage facilities of the cloud-based collaboration platform and may include any one

data storage components operating at any one or more physical devices. For example, the data repository may be distributed across multiple network connected machines. As mentioned, the data repository may store digital content that is hosted in the cloud-based collaboration platform (i.e., content uploaded by users/collaborators). The data repository may also store code including the instructions, rules, models, etc. for performing one or more processes of the cloud-based collaboration platform. For example, the cloud-based collaboration platform may host various user-facing applications (e.g., Microsoft® Word) that can be used by users/collaborators for interacting with stored digital content. Similarly, the cloud-based collaboration platform may host various processing applications that operate in the background, for example, to perform various automated tasks.

**[0027]** A content item (F1-Fn) stored/accessed within the cloud-based collaboration platform can generally include any type of digital or electronic content that can be viewed or accessed via an electronic device. The digital content can include, for example, pdf files, docs, slides (e.g., PowerPoint slides), images, audio files, multimedia content, web pages, blogs, etc. (which are collectively referred to herein illustratively and without limitation as a “file”).

**[0028]** A workspace can generally refer to any grouping of a set of digital content in the collaboration platform. The grouping can be created, identified, or specified by a user or through other means. This user may be a creator user or administrative user, for example. In general, a workspace can be associated with a set of users or collaborators which have access to the content included therein. The levels of access (e.g., based on permissions or rules) of each user or collaborator to access the content in a given workspace may be the same or may vary among the users. Each user may have their own set of access rights to every piece of content in the workspace, or each user may be different access rights to different pieces of content. Access rights may be specified by a user associated with a work space and/or a user who created/uploaded a particular piece of content to the workspace, or any other designated user or collaborator.

**[0029]** In general, the collaboration platform allows multiple users or collaborators to access or collaborate efforts on content items such each user can upload, open, view/preview/see, edit, revise, comment upon, and/or annotate specific content items through their own user devices. For example, a user can upload a document to a work space for other users to access (e.g., for viewing, editing, commenting, signing-off, or otherwise manipulating). The user can login to the online platform and upload the document (or any other type of content item) to an existing work space or to a new work space. The document can be shared with existing users or collaborators in a work space.

**[0030]** In some embodiments, anomaly detection is performed by using a probability model that is based upon clusters, where the clusters are formed from entity/object relationships specifically identified or correlated to the interaction/storage aspects of the cloud-based collaboration platform. In particular, anomaly detection is performed with regards to user-file interactions using a probability model that predicts the likelihood of a user interacting with a file. In this approach, low probability events that satisfy certain other conditions are deemed to be suspicious and worthy of further analysis.

[0031] FIG. 2 shows a flowchart of an approach to use the system of FIG. 1 to perform anomaly detection. At 202, the process collects certain information that are used to generate the probability model 116, including user information 106, interaction/activity information 108, and file/content information 110. The user information 106 pertains to data regarding user/collaborators in the system, such as user identifiers, user roles, permissions, as well as other relevant data about users within the system. The interaction/activity information 108 pertains to any action that is performed on a content item by a given user, such as form example, upload, download, edit, change, modification, comment, view, and/or open of a file within the collaboration platform. The file/content information 110 pertains to information about files/content within the system, such as the file identifier, owner, edit times, and any other relevant information about the file.

[0032] At step 204, a model generator 112 generates a probability model 116 that is based upon clustering of user-file activity graph(s). The model generator 112 generates the model 116 based upon information received from the collaboration platform, including the user information 106, the interaction/activity information 108, and the file/content information 110. Further details regarding an illustrative approach to generate probability model 116 is describe below in conjunction with FIGS. 3-8.

[0033] At 206, an activity stream 114 is accessed by the anomaly-detection/behavior-analysis system 120. The activity stream 114 corresponds to a sequence of activities and interactions that take place between users and files within the collaboration system. Such activities and interactions include, for example, upload, downloads, comments, views/previews, and edits of files within the collaboration system.

[0034] Scoring is then performed, at 208, upon the activities that are identified within the activity stream. The scoring is performed using the probability model 116 to generate scores for some or all of the activities. Anomalies 122 can then be identified from the scored activities. Further details regarding an illustrative approach to perform scoring is described below in conjunction with FIGS. 9-10.

[0035] Thereafter, at 210, the anomalous events 122 can be reported as possible threats within the system. The reporting can be made, for example, to administrators at one or more admin consoles 130.

[0036] In some embodiments, the probability model can be embodied as one or more files that include the model contents described herein, which therefore results in a new type of file that is usable to perform behavior-based detection of potentially hostile, harmful or anomalous activity operations in the computing system, where the instant model-based approach embodied in the new type of file differs from traditional approaches to anomaly detection. As discussed in more detail below, the approach described herein therefore improves the functionality of the computing system by being able to identify potential threats in the computing system based upon paradigms of user-file interactions that exist in cloud-based management systems. In addition, the approach described herein improves the functionality of the computing system by allowing very efficient use of computing resources (reduced memory usage and faster processing) to perform the anomaly detection.

[0037] This disclosure will now describe an illustrative approach for generating a probability model according to some embodiments of the invention, where the probability

model predicts the likelihood of a user interacting with a file. In this approach, low probability events that satisfy certain other conditions are deemed to be suspicious and worthy of further analysis.

[0038] FIG. 3 shows a flowchart of an approach to generate a probability model according to some embodiments of the invention. At 302, one or more user-to-file activity graph(s) are generated. The activity graph is formed by identifying the interactions that occur between specific users and specific files within the system.

[0039] FIG. 4 illustrates an example user-to-file activity graph for users 1, 2, 3 with respect to their interactions with files F1, F2, and F3. The users and files correspond to nodes within the graph and the interactions correspond to edges within the graph. A “weight” is associated with each user-to-file interaction/edge in the activity graph. In some embodiments, the weight is calculated by summing the numbers of activities (and/or certain selected types of activities) between a specific user and a specific file. In the illustrated example of FIG. 4, it can be seen that certain of the edges are represented with thicker lines and certain edges have thinner lines. The thicker lines represent weights that indicate greater user-to-file activities (e.g.,  $W_{U1-F1}$ ,  $W_{U1-F2}$ ,  $W_{U2-F1}$ ,  $W_{U2-F2}$ , and  $W_{U3-F3}$ ), while the thinner lines indicate lower levels of user-to-file activities (e.g.,  $W_{U1-F3}$ ,  $W_{U2-F3}$ ,  $W_{U3-F1}$ , and  $W_{U3-F2}$ ).

[0040] The activity graph can be formed by, for every file, generating a list of all users connected to that file. In some embodiments, pruning can be performed to prune the files with very high fan out. For example, this can be accomplished by sorting the files by their fan-out, and removing the top k % (e.g., where k is in the order of 0.1). In addition, the weights can be normalized based on the file fan out. To perform normalization in some embodiments, assume that “F” is the file’s fan-out, e.g., the number of users connected to the file in the activity graph. Then, normalize every weight “w” of every user connected to this file by dividing it by an appropriate factor, e.g.,

$$w' = \frac{w}{1 + \alpha \log(f)}$$

where w is the original weight, w' is the normalized weight, and a is the normalization factor.

[0041] At 304, the user-to-file graph is collapsed to form a user-to-user file. This step is performed by factoring out the files in the file-user graph to generate the user-user graph. For example, this can be accomplished by, for every file shared by two users, adding the smaller file-user weight to the user-user link.

[0042] FIG. 5 illustrates the process to generate the user-to-user graph from the user-to-file graph. Here, graph 502 represents the user-to-file activity graph. This graph 502 is reproduced from the graph described above in FIG. 4. Next, user-to-user weights are calculated from the user-to-file graph. This is represented in graph 504, where for example,  $W_{U1-U2}$  represents the weights of the link between User 1 and User 2. Similarly,  $W_{U1-U3}$  represents the weights of the link between User 1 and User 3, and  $W_{U2-U3}$  represents the weights of the link between User 2 and User 3.

[0043] Any suitable approach to determining weight values can be employed in embodiments of the invention. In at least one embodiment, for every file shared by two users, the



smaller file-to-user weight is added to a specific user-to-user link. For example, an illustrative approach to implementing weight  $W_{U1-U2}$  is where this weight represents the smaller of the weights for the common files shared by User 1 and User 2. However, as would be understood by those of ordinary skill in the art, other approaches may also be employed to calculate these weights.

[0044] As can be seen in the figure, the weight  $W_{U1-U2}$  of the link between user 1 and user 2 is represented by a thick line, indicating that the smaller of the weights for the common files shared by these two users (F1 and F2) is significant, indicating that these two users both tend to access the same set of files. On the other hand, the weights  $W_{U1-U3}$  and  $W_{U2-U3}$  are represented by thin lines, indicating that the smaller of the weights for the common files shared by these users is very small, indicating that user 1 and user 2 tend not to access the same files as user 3. Finally, as shown in graph 506, the user-to-user graph is formed by collapsing graph 504 to include just the users, as well as the links between the users.

[0045] At step 306, clustering is performed to cluster the user-to-user graph. This is illustrated in FIG. 6, where clustering is applied to graph 506 to identify users 1 and 2 in revised graph 602 as belonging to a first cluster, while user 3 is identified as belong to a second cluster. In some embodiments, Markov clustering can be used as the technique to cluster the user-user graph.

[0046] Next, at step 308, files are assigned to each of the clusters. The files to user clusters based on “affinity”, where the affinity of the file is computed relative to each one of the user clusters. The file is assigned to the cluster that has highest affinity to it.

[0047] With regards to affinity, it is noted that every cluster is a set of users and files. In some embodiments, the affinity of a file to a cluster is the sum of the weights of the connections between the file and all the users in the cluster. FIG. 7 provides an illustration of this approach, where files F1 and F2 are assigned to cluster 1. This is because the sum of the high weights of the connections between files F1 and F2 to the users 1 and 2 in cluster 1 (the weights  $W_{U1-F1}$ ,  $W_{U1-F2}$ ,  $W_{U2-F1}$ ,  $W_{U2-F2}$  with the thick lines) far exceeds the sum of the weights of the connections between files F1 and F2 to user 3 in cluster 2 (the weights  $W_{U3-F1}$  and  $W_{U3-F2}$  with the thin lines). In contrast, file 3 is assigned to cluster 2. This is because the sum of the high weight of the connection between file F3 to user 3 in cluster 2 (the weight  $W_{U3-F3}$  with the thick line) far exceeds the sum of the weights of the connections between file F3 to users 1 and 2 in cluster 1 (the weights  $W_{U1-F3}$  and  $W_{U2-F3}$  with the thin lines).

[0048] At step 310, feature vectors are created, including user-to-cluster vectors at step 312 and file-to-cluster vectors at step 314. FIG. 8 illustrates example feature vectors 804 between a user/file 802 and one or more clusters 806. Each feature vector 804 is associated with a vector affinity value/weight 808. The feature vector is an N-dimensional vector, where N is the number of clusters. For a file, the value at position k is the affinity of the file to cluster number k. For a user, the value at position k is the affinity of the user to cluster number k.

[0049] In some embodiments, for every file F and every user U, the approach computes feature vectors  $V(F)$  and  $W(U)$ . In some embodiments, V and W have the same dimensions. The basic idea is that, if the feature extraction

works well, the probability that U interacts with F is a monotonically increasing function of  $W(U) \cdot V(F)$ , where  $W(U) \cdot V(F)$  is the inner product of  $W(U)$  and  $V(F)$ .

[0050] The starting point for extracting the file and user feature vectors is the activity graph generated in the previous steps. The steps to calculate the feature vectors using activity graphs include some or all of (1): Cluster the activity graph into n clusters as discussed above, where these clusters are  $c_1, c_2, \dots, c_n$ ; (2) For the file F, calculate the affinity of F to  $c_1, c_2, \dots, c_n$ , resulting in  $a_1, a_2, \dots, a_n$ , where  $V(F) = (a_1, a_2, \dots, a_n)$ ; (3) For the user U, calculate the affinity of U to  $c_1, c_2, \dots, c_n$ , resulting in  $b_1, b_2, \dots, b_n$ , where  $W(U) = (b_1, b_2, \dots, b_n)$ .

[0051] With regards to affinity, it is noted that every cluster is a set of users and files. The affinity of a file to a cluster is the sum of the weights of the connections between the file and all the users in the cluster. The affinity of a user to a cluster is the sum of the weights of the connections between the user and all the files in the cluster. The weights that are used for computing affinity can be the original, non-normalized weights.

[0052] In some embodiments, the feature vector integrates all activities for a respective user/file. For example, this means that all activities for a user relative to a cluster are combined together for a combined user-cluster vector. In an alternative embodiment, the clusters can be established having smaller levels of granularity. For example, for a given user, each type of interaction by that user to a given cluster results in a separate feature cluster, where a first vector pertains to all “edit” actions by that user pertaining to a cluster, a second vector pertains to all “upload” actions by that user for the cluster, a third vector pertains to all “view” or “open” actions by that user for the cluster, a fourth vector pertains to all “comment” actions by that user for the cluster, and so on. Additional dimensions can be considered to create even finer granularities of vectors, e.g., to account for time of day, type of content, file extensions, etc.

[0053] Once the feature vectors have been generated, they can then be used for user behavioral analysis. FIG. 9 shows a flowchart of an approach to generate scoring to perform behavioral analysis according to some embodiments of the invention.

[0054] At step 902, interactions between users and files are identified, e.g., from an activity stream. At 904, a score is generated for each of the identified interactions. The score can be determined, for example, by using the dot product/inner product of the user and file feature vectors. This is illustrated in FIG. 10, where the scoring is determined based upon the user and file vectors for User 1 and File 1 that results in a first score, while the scoring between the user and file vectors for User 3 and File 1 results in a second score. It is noted that any suitable approach for scoring can be employed in embodiments of the invention, and thus the inventive concepts disclosed herein are not limited to specific scoring approaches such as dot product/inner product or cosine similarity, unless explicitly claimed as such.

[0055] In this illustrative example on the left side of the figure, it can be seen that the weight of the user and file vectors for User 1 and File 1 all correspond to thick lines indicating high vector weights. This shows a historical access pattern of consistent accesses to File 1 by User 1. Therefore, the dot product of the vectors in this situation will likely results in a very high score.

[0056] In contrast, on the right side of the figure, it can be seen that the weight of the user vectors for User 3 correspond to thin lines representing low vector weights, which are markedly different from the thick lines for the file vectors representing high vector weights. This disparity shows a historical access pattern of inconsistent accesses to File 1 by User 3. Therefore, the dot product of the vectors in this situation will likely results in a very low score.

[0057] As previously noted, the feature vectors may be generated at various levels of granularity. This means that scoring may be determined as well at the different levels of granularity, in correspondence with the type of feature vector that is used.

[0058] At step 906, the interactions are then sorted by their scores. Next, at step 908, the top or bottom set of interactions are identified. The choice of whether to take the top or bottom scores is dependent upon the type of scoring that is used. For example, if dot product scoring is employed, then the bottom nth interactions are identified, e.g., the bottom 1% of the scored interactions. One reason for taking the top (or bottom) percentage is to reduce the amount of data that would otherwise need to be processed to identify the possible threats. By significantly reducing the quantity of data to be processed in downstream processing (by eliminating the vast majority of the scored data), this approach eliminates interactions that are less likely to be a threat, reduces the amount of data that needs to be held in memory for processing, and increases the speed of processing since less data would therefore need to be operated upon to identify the potential threats.

[0059] At 910, the possible threats are then identified from the data set. In some embodiments, this is accomplished by determining the count of users that appear in the bottom 1% of the scored interactions, where the user(s) having the greatest number of interactions in the bottom 1% are highlighted as behavioral anomalies.

[0060] In some embodiments, the possible threats are identified by creating an alert that is sent to an administrator. In an alternate embodiment, the potential threat is automatically processed to mitigate possible exposure to the threat, e.g., by blocking and/or throttling access to one or more files by one or more users. For example, one possibility is to temporarily block all file access in the cloud management system by users identified in the bottom nth interactions. In an alternate embodiment, access is only blocked for the identified users for interactions that have scores below (or above) a certain threshold level that are indicative of anomalous access. In yet another embodiment, throttling is applied to limit the amount and/or rate of content access that is accessible for identified users/files.

[0061] Therefore, what has been described is an improved approach to implement anomaly detection, where the approach uses the probabilities of user-file interaction (determined using the inner product of user and file feature vectors) to identify sequence(s) of low-probability events that may be anomalous. Here, a single low-probability event is when a file that is normally accessed only by a particular cluster is accessed by a user of a different cluster who does not access other clusters often. The file access patterns and user access patterns are thus both employed to identify user behavior that are out-of-character.

[0062] It is noted that the above-described cluster-based approach to model generation is advantageous because it allows the effective capture of inter-dependencies in the data

that correlates to how people actually work together within the enterprise/system. This is because the clusters correspond to team activities and reflect the natural organization of an enterprise. In addition, the cluster-based approach provides a flexible representation having the ability to incorporate different node and edge types within the representations, while making the graphs less sparse. In effect, the models based upon clustering allows grouping of documents based upon patterns of use and are highly predictive of user behavior. Moreover, the present approach is also very robust to adversaries, making it difficult for fraudsters to understand the whole network and how specific entities/interactions fit in.

[0063] The inventive approach described herein is also advantageous over conventional machine learning techniques. For example, many conventional supervised machine learning techniques are dependent upon the existence of labelled data, and will therefore fail if accurate labels are not available for collected data. In contrast, the present approach is highly effective even if labeled data does not exist, since the inventive technique does not need to ground truth labels for what kind of activity is anomalous or not. In addition, conventional machine learning techniques often fail in large systems having extensively changing interactions that are very complex. In contrast, the inventive approach is expressly adapted to handle large scale and complex interactions, where the complex interactions between users and files may change over time. In addition, many conventional techniques for anomaly detection rely upon sniffing of network traffic to gather packets that are analyzed based upon packet metadata, where this approach is suboptimal because it is difficult if not impossible to always identify with certainty specific user/file interactions just from looking at individual packets (such as distinguishing open, edit, comment interactions from one another). In contrast, the inventive approach works upon an activity stream and directly performs anomaly detection using user-file interactions that exist in content management systems.

## System Architecture Overview

### Additional System Architecture Examples

[0064] FIG. 11A depicts a block diagram of an instance of a computer system 10A00 suitable for implementing embodiments of the present disclosure. Computer system 10A00 includes a bus 1006 or other communication mechanism for communicating information. The bus interconnects subsystems and devices such as a CPU, or a multi-core CPU (e.g., processor 1007), a system memory (e.g., main memory 1008, or an area of random-access memory RAM), a non-volatile storage device or area (e.g., ROM 1009), an internal or external storage device 1010 (e.g., magnetic or optical), a data interface 1033, a communications interface 1014 (e.g., PHY, MAC, Ethernet interface, modem, etc.). The aforementioned components are shown within processing element partition 1001, however other partitions are possible. The shown computer system 10A00 further comprises a display 1011 (e.g., CRT or LCD), various input devices 1012 (e.g., keyboard, cursor control), and an external data repository 1031.

[0065] According to an embodiment of the disclosure, computer system 10A00 performs specific operations by processor 1007 executing one or more sequences of one or more program code instructions contained in a memory.

Such instructions (e.g., program instructions **10021**, program instructions **10022**, program instructions **10023**, etc.) can be contained in or can be read into a storage location or memory from any computer readable/usable medium such as a static storage device or a disk drive. The sequences can be organized to be accessed by one or more processing entities configured to execute a single process or configured to execute multiple concurrent processes to perform work. A processing entity can be hardware-based (e.g., involving one or more cores) or software-based, and/or can be formed using a combination of hardware and software that implements logic, and/or can carry out computations and/or processing steps using one or more processes and/or one or more tasks and/or one or more threads or any combination therefrom.

[**0066**] According to an embodiment of the disclosure, computer system **10A00** performs specific networking operations using one or more instances of communications interface **1014**. Instances of the communications interface **1014** may comprise one or more networking ports that are configurable (e.g., pertaining to speed, protocol, physical layer characteristics, media access characteristics, etc.) and any particular instance of the communications interface **1014** or port thereto can be configured differently from any other particular instance. Portions of a communication protocol can be carried out in whole or in part by any instance of the communications interface **1014**, and data (e.g., packets, data structures, bit fields, etc.) can be positioned in storage locations within communications interface **1014**, or within system memory, and such data can be accessed (e.g., using random access addressing, or using direct memory access DMA, etc.) by devices such as processor **1007**.

[**0067**] The communications link **1015** can be configured to transmit (e.g., send, receive, signal, etc.) communications packets **1038** comprising any organization of data items. The data items can comprise a payload data area **1037**, a destination address **1036** (e.g., a destination IP address), a source address **1035** (e.g., a source IP address), and can include various encodings or formatting of bit fields to populate the shown packet characteristics **1034**. In some cases, the packet characteristics include a version identifier, a packet or payload length, a traffic class, a flow label, etc. In some cases, the payload data area **1037** comprises a data structure that is encoded and/or formatted to fit into byte or word boundaries of the packet.

[**0068**] In some embodiments, hard-wired circuitry may be used in place of or in combination with software instructions to implement aspects of the disclosure. Thus, embodiments of the disclosure are not limited to any specific combination of hardware circuitry and/or software. In embodiments, the term “logic” shall mean any combination of software or hardware that is used to implement all or part of the disclosure.

[**0069**] The term “computer readable medium” or “computer usable medium” as used herein refers to any medium that participates in providing instructions to processor **1007** for execution. Such a medium may take many forms including, but not limited to, non-volatile media and volatile media. Non-volatile media includes, for example, optical or magnetic disks such as disk drives or tape drives. Volatile media includes dynamic memory such as a random-access memory.

[**0070**] Common forms of computer readable media include, for example, floppy disk, flexible disk, hard disk,

magnetic tape, or any other magnetic medium; CD-ROM or any other optical medium; punch cards, paper tape, or any other physical medium with patterns of holes; RAM, PROM, EPROM, FLASH-EPROM, or any other memory chip or cartridge, or any other non-transitory computer readable medium. Such data can be stored, for example, in any form of external data repository **1031**, which in turn can be formatted into any one or more storage areas, and which can comprise parameterized storage **1039** accessible by a key (e.g., filename, table name, block address, offset address, etc.).

[**0071**] Execution of the sequences of instructions to practice certain embodiments of the disclosure are performed by a single instance of the computer system **10A00**. According to certain embodiments of the disclosure, two or more instances of computer system **10A00** coupled by a communications link **1015** (e.g., LAN, PTSN, or wireless network) may perform the sequence of instructions required to practice embodiments of the disclosure using two or more instances of components of computer system **10A00**.

[**0072**] The computer system **10A00** may transmit and receive messages such as data and/or instructions organized into a data structure (e.g., communications packets **1038**). The data structure can include program instructions (e.g., application code **1003**), communicated through communications link **1015** and communications interface **1014**. Received program code may be executed by processor **1007** as it is received and/or stored in the shown storage device or in or upon any other non-volatile storage for later execution. Computer system **10A00** may communicate through a data interface **1033** to a database **1032** on an external data repository **1031**. Data items in a database can be accessed using a primary key (e.g., a relational database primary key).

[**0073**] The processing element partition **1001** is merely one sample partition. Other partitions can include multiple data processors, and/or multiple communications interfaces, and/or multiple storage devices, etc. within a partition. For example, a partition can bound a multi-core processor (e.g., possibly including embedded or co-located memory), or a partition can bound a computing cluster having plurality of computing elements, any of which computing elements are connected directly or indirectly to a communications link. A first partition can be configured to communicate to a second partition. A particular first partition and particular second partition can be congruent (e.g., in a processing element array) or can be different (e.g., comprising disjoint sets of components).

[**0074**] A module as used herein can be implemented using any mix of any portions of the system memory and any extent of hard-wired circuitry including hard-wired circuitry embodied as a processor **1007**. Some embodiments include one or more special-purpose hardware components (e.g., power control, logic, sensors, transducers, etc.). A module may include one or more state machines and/or combinational logic used to implement or facilitate the performance characteristics of systems for dynamically monitoring code execution activity to identify and manage inactive code.

[**0075**] Various implementations of the database **1032** comprise storage media organized to hold a series of records or files such that individual records or files are accessed using a name or key (e.g., a primary key or a combination of keys and/or query clauses). Such files or records can be organized into one or more data structures (e.g., data structures used to implement or facilitate aspects of dynamically

monitoring code execution activity to identify and manage inactive code). Such files or records can be brought into and/or stored in volatile or non-volatile memory.

**[0076]** FIG. 11B depicts a block diagram of an instance of a cloud-based environment **10B00**. Such a cloud-based environment supports access to workspaces through the execution of workspace view code (e.g., workspace access code **10521** and workspace access code **10522**. Workspace access code can be executed on any of the shown user devices **1056** (e.g., laptop device **10564**, workstation device **10565**, IP phone device **10563**, tablet device **10562**, smart phone device **10561**, etc.), or on one or more processing elements. A group of users can form a collaborator group **1058**, and a collaborator group can be comprised of any types or roles of users. For example, and as shown, a collaborator group can comprise a user collaborator, an administrator collaborator, a creator collaborator, etc. Any user can use any one or more of the user devices, and such user devices can be operated concurrently to provide multiple concurrent sessions and/or other techniques to access workspaces through the workspace access code.

**[0077]** A portion of workspace access code can reside in and be executed on any user device. Also, a portion of the workspace access code can reside in and be executed on any computing platform, including in a middleware setting. As shown, a portion of the workspace access code resides in and can be executed on one or more processing elements (e.g., processing element **10531**). The workspace access code can interface with storage devices such the shown network storage **1055**. Storage of workspaces and/or any constituent files or objects, and/or any other code or scripts or data can be stored in any one or more storage partitions (e.g., storage partition **10541**). In some environments, a processing element includes forms of storage such as RAM and/or ROM and/or FLASH, and/or other forms of volatile and non-volatile storage.

**[0078]** A stored workspace can be populated via an upload (e.g., an upload from a user device to a processing element over an upload network path **1057**). A stored workspace can be delivered to a particular user and/or shared with other

particular users via a download (e.g., a download from a processing element to a user device over a download network path **1059**).

**[0079]** In the foregoing specification, the disclosure has been described with reference to specific embodiments thereof. It will, however, be evident that various modifications and changes may be made thereto without departing from the broader spirit and scope of the disclosure. For example, the above-described process flows are described with reference to a particular ordering of process actions. However, the ordering of many of the described process actions may be changed without affecting the scope or operation of the disclosure. The specification and drawings to be regarded in an illustrative sense rather than in a restrictive sense.

What is claimed is:

1. A method of predicting user-file interactions comprising:
  - determining a model that represents interactions between multiple users and multiple shared files that are commonly accessible among the multiple users in a collaboration environment at least by converting a user-to-file graph into a converted graph, wherein
    - at least a file node of multiple file nodes is removed from the user-to-file graph in determining the model, and
    - the user-to-file graph comprises multiple user nodes respectively connected to the multiple file nodes respectively representing the multiple shared files with one or more corresponding edges that represent the interactions between the multiple user nodes and the multiple file nodes;
  - generating one or more clusters for the model from the converted graph;
  - receiving interactions between at least some of the multiple users and at least some of the multiple shared files in the collaboration system; and
  - determining respective scores for the interactions using the model.

\* \* \* \* \*