

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent Application Publication

20250267185

Kind Code

A1

Publication Date

August 21, 2025

Inventor(s)

Kommula; Raja et al.

SELF-LEARNING SERVICE SCHEDULER FOR SMART NICS

Abstract

An example method comprises determining, by an edge services controller, based on a respective predicted resource utilization value for each of a plurality of servers, a corresponding server weight for each of the plurality of servers; the plurality of servers comprising respective network interface cards (NICs), wherein each NIC of the plurality of NICs comprises an embedded switch and a processing unit coupled to the embedded switch; determining, by the edge services controller, based on a respective predicted resource utilization value for each of a plurality of services, a corresponding application weight for each of the plurality of services; and scheduling, by the edge services controller, based on the respective server weight for a server of the plurality of servers and the respective application weight for the service, a service of the plurality of services on the server.

Inventors: Kommula; Raja (Cupertino, CA), Sunkada; Ganesh Byagoti Matad (Bengaluru, IN), Sridhar; Thayumanavan (Sunnyvale, CA), Krishnamoorthy; Rajasree (Fremont, CA), Yavatkar; Raj (Los Gatos, CA), Gupta; Jit (Philadelphia, PA), Kant; Krishna (Philadelphia, PA)

Applicant: Juniper Networks, Inc. (Sunnyvale, CA)

Family ID: 1000008589318

Appl. No.: 19/189729

Filed: April 25, 2025

Related U.S. Application Data

parent US continuation 18640970 20240419 parent-grant-document US 12289364 child US 19189729

parent US continuation 18064803 20221212 parent-grant-document US 11968251 child US 18640970

Publication Classification

Int. Cl.: H04L67/1008 (20220101); H04L41/16 (20220101)

U.S. Cl.:

CPC H04L67/1008 (20130101); H04L41/16 (20130101);

Background/Summary

[0001] This application is a continuation of U.S. application Ser. No. 18/640,970, filed Apr. 19, 2024, which is a continuation of U.S. patent application Ser. No. 18/064,803, filed Dec. 12, 2022, now issued U.S. Pat. No. 11,968,251, the entire contents of which is incorporated herein by reference.

TECHNICAL FIELD

[0002] The disclosure relates to computer networks.

BACKGROUND

[0003] In a typical cloud-based data center environment, a large collection of interconnected servers provide computing and/or storage capacity to run various applications. For example, a data center may comprise a facility that hosts applications and services for subscribers, i.e., customers of a data center provider. The data center may, for example, host all of the infrastructure equipment, such as networking and storage systems, redundant power supplies, and environmental controls. In a typical data center, clusters of storage servers and application servers (compute nodes) are interconnected via a high-speed switch fabric provided by one or more tiers of physical network switches and routers. More sophisticated data centers provide infrastructure spread throughout the world with subscriber support equipment located in various physical hosting facilities.

[0004] Connectivity between the server and the switch fabric is provided by a hardware module called a Network Interface Card (NIC). A conventional NIC includes an application-specific integrated circuit (ASIC) to perform packet forwarding, which includes some basic Layer 2/Layer 3 (L2/L3) functionality. In conventional NICs, the packet processing, policing and other advanced functionality, known as the “datapath,” is performed by the host CPU, i.e., the CPU of the server that includes the NIC. However, some NIC vendors may include an additional processing unit in the NIC itself to offload at least some of the datapath processing from the host CPU to the NIC. The processing unit in the NIC may be, e.g., a multi-core ARM processor with some hardware acceleration provided by a Data Processing Unit (DPU), Field Programmable Gate Array (FPGA), and/or an ASIC. A processing unit may be alternatively referred to as a DPU. NICs that include such augmented datapath processing capabilities are typically referred to as SmartNICs.

SUMMARY

[0005] In general, techniques are described for an edge services platform that leverages processing units of NICs to augment the processing and networking functionality of a network of servers that include the NICs. Features provided by the edge services platform may include, e.g., orchestration of NICs; API driven deployment of services on NICs; NIC addition, deletion, and replacement; monitoring of services and other resources on NICs; and management of connectivity between various services running on the NICs. More specifically, this disclosure describes an edge services platform that implements a self-learning scheduler to determine placements of service instances based on predicted future application/service resource utilizations and server resource utilizations. The self-learning scheduler learns about resource utilization requirements of one or more services, and resource utilization patterns of hardware resources of one or more servers, including resource

utilization patterns of processing units of NICs of the servers. In some examples, the edge services platform uses a machine learning model to inform scheduling of a service of the one or more services to a server of the one or more servers. The machine learning model is trained by analyzing a first set of historical utilization data for the one or more servers and a second set of historical utilization data for the one or more services. The trained machine learning model predicts resource utilization values for the one or more servers and the one or more services.

[0006] Services running in a typical enterprise data center environment may follow a predictable pattern of server compute and network resource usage. For example, administrators may schedule backup services during the overnight hours. Payroll processing software may run on a weekly, bimonthly, or monthly schedule. Branch offices may upload daily transaction data to the central office at the end of the day. A typical use case is a backup service that follows a pattern of a large number of elephant flows from servers to storage devices during the backup process. Backup software may use jumbo packets to minimize the total number of required transactions. Administrators may pre-configure a virtual local area network (VLAN) for backup traffic.

[0007] When services follow predictable resource usage patterns, these services cause predictable resource usage patterns on the servers running the services. These patterns can be used to schedule the services onto appropriate servers to utilize server resources efficiently, and to maximize the performance of the services. Traditional schedulers make scheduling decisions based on the availability of resources on each of a plurality of servers at the time of scheduling, along with one or more user-configured static scheduling policies such as server affinity and server taints.

[0008] Since traditional schedulers are not aware of future utilization trends for services and resources on the plurality of servers, high resource utilization services may be scheduled onto a server which then faces a resource deficiency subsequent to the scheduling. Likewise, one or more low resource utilization services may be scheduled onto a server for which additional resources may become available subsequent to the scheduling. These scheduling inefficiencies may result in degraded service performance or underutilization of server resources.

[0009] The techniques may provide one or more technical advantages that realize one or more practical applications. For example, in contrast to traditional schedulers, the self-learning scheduler described herein acquires knowledge about future resource requirements of services and resource availabilities of servers, including the processing units of NICs of such servers. The self-learning scheduler uses a machine learning model to make predictions of future resource requirements based on usage telemetry data relative to start times for each of the services. The self-learning scheduler uses the machine learning model to make predictions of future resource availabilities for servers based on resource usage telemetry of service instances. Based on the predicted future resource requirements and predicted future resource availabilities, the self-learning scheduler schedules a service having a high future resource requirement with a server having a high future resource availability, such that a service requiring high resources is scheduled for deployment on a server having more resources. In this way, the techniques may improve overall service performance across multiple service instances deployed to multiple servers, reduce performance bottlenecks from oversubscribed servers, and/or facilitate enhanced reliability and performance of service instances by ensuring sufficient resources for future deployments.

[0010] In one example, this disclosure describes a system comprising a plurality of servers comprising respective network interface cards (NICs), wherein each NIC of the plurality of NICs comprises an embedded switch and a processing unit coupled to the embedded switch; and an edge services controller configured to determine, based on a corresponding predicted resource utilization value for the processing unit of the corresponding NIC of each of the plurality of servers, a corresponding server weight for each of the plurality of servers; determine, based on a corresponding predicted resource utilization value for each of the plurality of services, a corresponding application weight for each of a plurality of services; and schedule a service of the plurality of services on a processing unit of the corresponding NIC of a server of the plurality of

servers based on the corresponding server weight for the server and the corresponding application weight for the service.

[0011] In another example, this disclosure describes a method comprising determining, by an edge services controller, based on a respective predicted resource utilization value for each of a plurality of servers, a corresponding server weight for each of the plurality of servers; the plurality of servers comprising respective network interface cards (NICs), wherein each NIC of the plurality of NICs comprises an embedded switch and a processing unit coupled to the embedded switch; determining, by the edge services controller, based on a respective predicted resource utilization value for each of a plurality of services, a corresponding application weight for each of the plurality of services; and scheduling, by the edge services controller, based on the respective server weight for a server of the plurality of servers and the respective application weight for the service, a service of the plurality of services on the server.

[0012] In another example, this disclosure describes a non-transitory computer-readable storage medium comprising instructions that, when executed, configure processing circuitry of a computing system to perform operations comprising: determining, by an edge services controller, based on a respective predicted resource utilization value for each of a plurality of servers, a corresponding server weight for each of the plurality of servers; the plurality of servers comprising respective network interface cards (NICs), wherein each NIC of the plurality of NICs comprises an embedded switch and a processing unit coupled to the embedded switch; determining, by the edge services controller, based on a respective predicted resource utilization value for each of a plurality of services, a corresponding application weight for each of the plurality of services; and scheduling, by the edge services controller, based on the respective server weight for a server of the plurality of servers and the respective application weight for the service, a service of the plurality of services on the server.

[0013] The details of one or more embodiments of this disclosure are set forth in the accompanying drawings and the description below. Other features, objects, and advantages will be apparent from the description and drawings, and from the claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

[0014] FIG. 1 is a block diagram illustrating an example network system having a data center in which examples of the techniques described herein may be implemented.

[0015] FIG. 2 is a block diagram illustrating an example computing device that uses a network interface card having a separate processing unit, to perform services managed by an edge services platform according to techniques described herein.

[0016] FIG. 3 is a conceptual diagram illustrating a data center with servers that each include a network interface card having a separate processing unit, controlled by an edge services platform, according to techniques of this disclosure.

[0017] FIG. 4 is a block diagram illustrating an example computing device that uses a network interface card having a separate processing unit, to perform services managed by an edge services platform according to techniques described herein.

[0018] FIG. 5 is a block diagram illustrating an example system for scheduling one or more services to one or more servers, according to techniques of this disclosure.

[0019] FIG. 6 is a data flow diagram illustrating an example method for scheduling one or more services to one or more servers, according to techniques of this disclosure.

[0020] FIG. 7 is a block diagram illustrating an example system for scheduling one or more services to one or more servers, according to techniques of this disclosure.

[0021] FIG. 8 is a data flow diagram illustrating an example method for scheduling one or more

services to one or more servers, according to techniques of this disclosure.

[0022] FIG. **9** is a flowchart for an example method performed by an edge services controller according to techniques of this disclosure.

[0023] FIG. **10** is a table showing an illustrative resource utilization pattern on a data processing unit.

[0024] FIG. **11** is a table showing a server and application weight matrix for multiple data processing units of servers, at each of a plurality of different time stamps.

[0025] Like reference characters denote like elements throughout the description and figures.

DETAILED DESCRIPTION

[0026] FIG. **1** is a block diagram illustrating an example network system **8** having a data center **10** in which examples of the techniques described herein may be implemented. In general, data center **10** provides an operating environment for applications and services for customer sites **11** having one or more customer networks coupled to data center **10** by a service provider network **7**. Data center **10** may, for example, host infrastructure equipment, such as networking and storage systems, redundant power supplies, and environmental controls. Service provider network **7** is coupled to a public network **4**. Public network **4** may represent one or more networks administered by other providers and may thus form part of a large-scale public network infrastructure, e.g., the Internet. For instance, public network **4** may represent a local area network (LAN), a wide area network (WAN), the Internet, a virtual LAN (VLAN), an enterprise LAN, a layer **3** virtual private network (VPN), an Internet Protocol (IP) intranet operated by the service provider that operates service provider network **7**, an enterprise IP network, or some combination thereof.

[0027] Although customer sites **11** and public network **4** are illustrated and described primarily as edge networks of service provider network **7**, in some examples, one or more of customer sites **11** and public network **4** are tenant networks within data center **10** or another data center. For example, data center **10** may host multiple tenants (customers) each associated with one or more virtual private networks (VPNs). Each of the VPNs may implement one of customer sites **11**.

[0028] Service provider network **7** offers packet-based connectivity to attached customer sites **11**, data center **10**, and public network **4**. Service provider network **7** may represent a network that is operated (and potentially owned) by a service provider to interconnect a plurality of networks. Service provider network **7** may implement Multi-Protocol Label Switching (MPLS) forwarding and, in such instances, may be referred to as an MPLS network or MPLS backbone. In some instances, service provider network **7** represents a plurality of interconnected autonomous systems, such as the Internet, that offers services from one or more service providers.

[0029] In some examples, data center **10** may represent one of many geographically distributed network data centers. As illustrated in the example of FIG. **1**, data center **10** may be a facility that provides network services for customers. A customer of the service provider may be a collective entity such as enterprises and governments or individuals. For example, a network data center may host web services for several enterprises and end users. Other exemplary services may include data storage, virtual private networks, traffic engineering, file service, data mining, scientific- or super-computing, and so on. Although illustrated as a separate edge network of service provider network **7**, elements of data center **10** such as one or more physical network functions (PNFs) or virtualized network functions (VNFs) may be included within the service provider network **7** core.

[0030] In this example, data center **10** includes storage and/or compute servers interconnected via switch fabric **14** provided by one or more tiers of physical network switches and routers, with servers **12A-12X** (herein, “servers **12**”) depicted as coupled to top-of-rack (TOR) switches **16A-16N**. This disclosure may refer to TOR switches **16A-16N** collectively, as “TOR switches **16**.” TOR switches **16** may be network devices that provide layer **2** (MAC) and/or layer **3** (e.g., IP) routing and/or switching functionality.

[0031] Servers **12** may also be referred to herein as “hosts” or “host devices.” Data center **10** may include many additional servers coupled to other TOR switches **16** of the data center **10**.

[0032] Switch fabric **14** in the illustrated example includes interconnected TOR switches **16** (or other “leaf” switches) coupled to a distribution layer of chassis switches **18A-18M** (collectively, “chassis switches **18**”). Chassis switches may also be referred to as “spine” or “core” switches. Although not shown in the example of FIG. **1**, data center **10** may also include one or more non-edge switches, routers, hubs, gateways, security devices such as firewalls, intrusion detection, and/or intrusion prevention devices, servers, computer terminals, laptops, printers, databases, wireless mobile devices such as cellular phones or personal digital assistants, wireless access points, bridges, cable modems, application accelerators, and/or other network devices.

[0033] In some examples, TOR switches **16** and chassis switches **18** provide servers **12** with redundant (e.g., multi-homed) connectivity to IP fabric **20** and service provider network **7**. Chassis switches **18** aggregate traffic flows and provide connectivity between TOR switches **16**. TOR switches **16** and chassis switches **18** may each include one or more processors and a memory and can execute one or more software processes. Chassis switches **18** are coupled to IP fabric **20**, which may perform layer **3** routing to route network traffic between data center **10** and customer sites **11** via service provider network **7**. The switching architecture of data center **10** shown in FIG. **1** is merely an example. Other switching architectures may have more or fewer switching layers, for instance. TOR switches **16** and chassis switches **18** may each include physical network interfaces.

[0034] In this disclosure, the terms “packet flow,” “traffic flow,” or simply “flow” each refer to a set of packets originating from a particular source device or endpoint and sent to a particular destination device or endpoint. A single flow of packets may be identified by the 5-tuple: <source network address, destination network address, source port, destination port, protocol>, for example. This 5-tuple generally identifies a packet flow to which a received packet corresponds. An n-tuple refers to any n items drawn from the 5-tuple. For example, a 2-tuple for a packet may refer to the combination of <source network address, destination network address> or <source network address, source port> for the packet. The term “source port” refers to a transport layer (e.g., TCP/UDP) port. A “port” may refer to a physical network interface of a NIC.

[0035] Each of servers **12** may be a compute node, an application server, a storage server, or other type of server. For example, each of servers **12** may represent a computing device, such as an x86 processor-based server, configured to operate according to techniques described herein. Servers **12** may provide Network Function Virtualization Infrastructure (NFVI) for a Network Function Virtualization (NFV) architecture.

[0036] Servers **12** may host endpoints for one or more virtual networks that operate over the physical network represented in FIG. **1** by IP fabric **20** and switch fabric **14**. Endpoints may include, e.g., virtual machines, containerized applications, or applications executing natively on the operating system or bare metal. Although described primarily with respect to a data center-based switching network, other physical networks, such as service provider network **7**, may underlay the one or more virtual networks.

[0037] Each of servers **12** includes at least one network interface card (NIC) of NICs **13A-13X** (collectively, “NICs **13**”). For example, server **12A** includes NIC **13A**. Each of NICs **13** includes at least one port. Each of NICs **13** may send and receive packets over one or more communication links coupled to the ports of the NIC.

[0038] In some examples, each of NICs **13** provides one or more virtual hardware components for virtualized input/output (I/O). A virtual hardware component for virtualized I/O may be a virtualization of a physical NIC **13** (the “physical function”). For example, in Single Root I/O Virtualization (SR-IOV), which is described in the Peripheral Component Interface Special Interest Group SR-IOV specification, the Peripheral Component Interface (PCI) express (PCIe) Physical Function of the network interface card (or “network adapter”) is virtualized to present one or more virtual network interface cards as “virtual functions” for use by respective endpoints executing on the server **12**. In this way, the virtual network endpoints may share the same PCIe physical hardware resources and the virtual functions are examples of virtual hardware components. As

another example, one or more servers **12** may implement Virtio, a para-virtualization framework available, e.g., for the Linux Operating System, that provides emulated NIC functionality as a type of virtual hardware component. As another example, one or more servers **12** may implement Open vSwitch to perform distributed virtual multilayer switching between one or more virtual NICs (vNICs) for hosted virtual machines, where such vNICs may also represent a type of virtual hardware component. In some instances, the virtual hardware components are virtual I/O (e.g., NIC) components. In some instances, the virtual hardware components are SR-IOV virtual functions and may provide SR-IOV with Data Plane Development Kit (DPDK)-based direct process user space access.

[0039] In some examples, one or more of NICs **13** include multiple ports. NICs **13** may be connected to one another via ports of NICs **13** and communications links to form a NIC fabric having a NIC fabric topology. Such a NIC fabric is the collection of NICs **13** connected to at least one other of NICs **13** and the communications links coupling NICs **13** to one another.

[0040] NICs **13A-13X** include corresponding processing units **25A-25X** (collectively, “processing units **25**”). Processing units **25** to offload aspects of the datapath from CPUs of servers **12**. One or more of processing units **25** may be a multi-core ARM processor with hardware acceleration provided by a Data Processing Unit (DPU), a Field Programmable Gate Array (FPGA), and/or an Application Specific Integrated Circuit (ASIC). Because NICs **13** include processing units **25**, NICs **13** may be referred to as “SmartNICs” or “GeniusNICs.”

[0041] In accordance with various aspects of the techniques of this disclosure, an edge services platform uses processing units **25** of NICs **13** to augment the processing and networking functionality of switch fabric **14** and/or servers **12** that include NICs **13**. In the example of FIG. **1**, network system **8** includes an edge services controller **28**. This disclosure may also refer to an edge services controller, such as edge services controller **28**, as an edge services platform controller.

[0042] Edge services controller **28** may manage the operations of the edge services platform within NIC **13s** in part by orchestrating services performed by processing units **25**; orchestrating API driven deployment of services on NICs **13**; orchestrating NIC **13** addition, deletion and replacement within the edge services platform; monitoring of services and other resources on NICs **13**; and/or management of connectivity between various services **133** running on the NICs **13**. In some examples, edge services controller **28** may include one or more computing devices, such as server devices, personal computers, intermediate network devices, or the like, configured to execute a distributed implementation of an edge services controller. In some examples, edge services controller **28** may be implemented using a single computing device.

[0043] Edge services controller **28** may communicate information describing services available on NICs **13**, a topology of a NIC fabric, or other information about the edge services platform to an orchestration system (not shown) or a controller **24**. Edge services controller **28** may be integrated within an overall controller **24** for computing infrastructure **8**. Example orchestration systems include OpenStack, vCenter by VMWARE, or System Center by MICROSOFT CORPORATION of Redmond, Washington. Example controllers include a controller for Contrail by JUNIPER NETWORKS or Tungsten Fabric. Controller **24** may be a network fabric manager. Additional information regarding a controller **24** operating in conjunction with other devices of data center **10** or other software-defined network is found in International Application Number PCT/US2013/044378, filed Jun. 5, 2013, and entitled “PHYSICAL PATH DETERMINATION FOR VIRTUAL NETWORK PACKET FLOWS;” and in U.S. Pat. No. 9,571,391, filed Mar. 26, 2014, and entitled “Tunneled Packet Aggregation for Virtual Networks,” each of which is incorporated by reference as if fully set forth herein.

[0044] Edge services controller **28** may be configured to determine a respective server weight for each of a plurality of servers **12**, based on a respective predicted resource utilization value for each of the plurality of servers **12**; determine a respective application weight for each of a plurality of services, based on a respective predicted resource utilization value for each of the plurality of

services; schedule a service of the plurality of services on a server of the plurality of servers **12** based on the respective server weight for the server and the respective application weight for the service.

[0045] In some embodiments, the edge services controller **28** may implement a self-learning scheduler that acquires knowledge about future resource requirements of services and resource availabilities of servers **12**. The self-learning scheduler may use a machine learning model to make predictions of future resource requirements based on usage telemetry data relative to start times for each of the services. The self-learning scheduler may use the machine learning model to make predictions of future resource availabilities for servers **12** based on resource usage telemetry of service instances. Based on the predicted future resource requirements and predicted future resource availabilities, the self-learning scheduler can push a service having a high future resource requirement to a server having a high future resource availability. For example, a service requiring high resources, such as a first instance of first service **29**, can be scheduled for deployment on NIC **13A** of server **12A**. Pursuant to this illustrative example, server **12A** may have a greater amount of future resource availability as compared to at least one other server **12X**.

[0046] FIG. **2** is a block diagram illustrating an example computing device **200** that uses a NIC **230** having a separate processing unit **25**, to perform services managed by an edge services platform according to techniques described herein. Computing device **200** of FIG. **2** may represent a real or virtual server and may represent an example instance of any of servers **12** of FIG. **1**. In the example of FIG. **2**, computing device **200** includes a bus **242** that couples hardware components of the hardware environment of computing device **200**. Specifically, in the example of FIG. **2**, bus **242** couples a Single Route Input/Output Virtualization (SR-IOV)-capable NIC **230**, a storage disk **246**, and a microprocessor **210**. In some examples, a front-side bus couples microprocessor **210** and memory device **244**. In some examples, bus **242** couples memory device **244**, microprocessor **210**, and NIC **230**. Bus **242** may represent a PCIe bus. In some examples, a direct memory access (DMA) controller may control DMA transfers among components coupled to bus **242**. In some examples, components coupled to bus **242** control DMA transfers among components coupled to bus **242**.

[0047] Microprocessor **210** may include one or more processors each including an independent execution unit (“processing core”) to perform instructions that conform to an instruction set architecture. Execution units may be implemented as separate integrated circuits (ICs) or may be combined within one or more multi-core processors (or “many-core” processors) that are each implemented using a single IC (i.e., a chip multiprocessor).

[0048] Disk **246** represents computer readable storage media that includes volatile and/or non-volatile, removable and/or non-removable media implemented in any method or technology for storage of information such as processor-readable instructions, data structures, program modules, or other data. Computer readable storage media includes, but is not limited to, random access memory (RAM), read-only memory (ROM), EEPROM, flash memory, CD-ROM, digital versatile discs (DVD) or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium that can be used to store the desired information and that can be accessed by microprocessor **210**.

[0049] Memory device **244** includes one or more computer-readable storage media, which may include random-access memory (RAM) such as various forms of dynamic RAM (DRAM), e.g., DDR2/DDR3 SDRAM, or static RAM (SRAM), flash memory, or any other form of fixed or removable storage medium that can be used to carry or store desired program code and program data in the form of instructions or data structures and that can be accessed by a computer. Memory device **244** provides a physical address space composed of addressable memory locations.

[0050] Network interface card (NIC) **230** includes one or more interfaces **232** configured to exchange packets using links of an underlying physical network. Interfaces **232** may include a port interface card having one or more network ports. NIC **230** also include an on-card memory **227** to,

e.g., store packet data. Direct memory access transfers between NIC **230** and other devices coupled to bus **242** may read/write from/to the memory **227**.

[0051] Memory device **244**, NIC **230**, disk **246**, and microprocessor **210** provide an operating environment for a software stack that executes a hypervisor **214** and one or more virtual machines **228** managed by hypervisor **214**. In general, a virtual machine provides a virtualized/guest operating system for executing applications in an isolated virtual environment. Because a virtual machine is virtualized from physical hardware of the host server, executing applications are isolated from both the hardware of the host and other virtual machines. Computing device **200** executes hypervisor **214** to manage virtual machines **228**. Example hypervisors include Kernel-based Virtual Machine (KVM) for the Linux kernel, Xen, ESXi available from VMWARE, Windows Hyper-V available from MICROSOFT, and other open-source and proprietary hypervisors. Hypervisor **214** may represent a virtual machine manager (VMM). Virtual machines **228** may host one or more applications, such as virtual network function instances. In some examples, a virtual machine **228** may host one or more VNF instances, where each of the VNF instances is configured to apply a network function to packets.

[0052] An alternative to virtual machines is the virtualized container, such as those provided by the open-source DOCKER Container application. Like a virtual machine, each container is virtualized and may remain isolated from the host machine and other containers. However, unlike a virtual machine, each container may omit an individual operating system and provide only an application suite and application-specific libraries. A container is executed by the host machine as an isolated user-space instance and may share an operating system and common libraries with other containers executing on the host machine. Thus, containers may require less processing power, storage, and network resources than virtual machines. As used herein, containers may also be referred to as virtualization engines, virtual private servers, silos, or jails. In some instances, the techniques described herein with respect to containers and virtual machines or other virtualization components.

[0053] While virtual network endpoints in FIG. 2 are illustrated and described with respect to virtual machines, other operating environments, such as containers (e.g., a DOCKER container) may implement virtual network endpoints. An operating system kernel (not shown in FIG. 2) may execute in kernel space and may include, for example, a Linux, Berkeley Software Distribution (BSD), another Unix-variant kernel, or a Windows server operating system kernel, available from MICROSOFT.

[0054] Hypervisor **214** includes a physical driver **225** to use a physical function provided by NIC **230**. In some cases, NIC **230** may also implement SR-IOV to enable sharing the physical network function (I/O) among virtual machines **224**. Each port of NIC **230** may be associated with a different physical function. The shared virtual devices, also known as virtual functions, provide dedicated resources such that each of virtual machines **228** (and corresponding guest operating systems) may access dedicated resources of NIC **230**, which therefore appears to each of virtual machines **224** as a dedicated NIC. Virtual functions may be lightweight PCIe functions that share physical resources with the physical function and with other virtual functions. NIC **230** may have thousands of available virtual functions according to the SR-IOV standard, but for I/O-intensive applications the number of configured virtual functions is typically much smaller.

[0055] Virtual machines **228** include respective virtual NICs **229** presented directly into the virtual machine **228** guest operating system, thereby offering direct communication between NIC **230** and virtual machines **228** via bus **242**, using the virtual function assigned for the virtual machine. This may reduce hypervisor **214** overhead involved with software-based, VIRTIO and/or vSwitch implementations in which a memory address space of hypervisor **214** within memory device **244** stores packet data and because copying packet data from NIC **230** to the memory address space of hypervisor **214** and from the memory address space of hypervisor **214** to memory address spaces of virtual machines **228** consumes cycles of microprocessor **210**. Microprocessor **210** is a server

resource and may be associated with a CPU utilization percentage. The CPU utilization percentage may be incorporated into the machine learning model.

[0056] NIC **230** may be associated with a network utilization for bandwidth to/from the ports of NIC **239**. The network utilization for bandwidth is a server resource and can be incorporated into the machine learning model. NIC **230** may further include a hardware-based Ethernet bridge **234**. Ethernet bridge **234** may be an example of an embedded switch **234**. Ethernet bridge **234** may perform layer 2 forwarding between virtual functions and physical functions of NIC **230**. Thus, in some cases, Ethernet bridge **234** provides hardware acceleration, via bus **242**, of inter-virtual machine **224** packet forwarding and hardware acceleration of packet forwarding between hypervisor **214** and any of virtual machines **224**. Hypervisor **214** may access the physical function via physical driver **225**. Ethernet bridge **234** may be physically separate from processing unit **25**. Processing unit **25** is a server resource and may be associated with a data processing unit (DPU) utilization percentage. The DPU utilization percentage can be incorporated into the machine learning model.

[0057] Computing device **200** may be coupled to a physical network switch fabric that includes an overlay network that extends a switch fabric from physical switches to software or “virtual” routers of physical servers coupled to the switch fabric, including virtual router **220**. Virtual routers may be processes or threads, or a component thereof, executed by the physical servers, e.g., servers **12** of FIG. **1**, that dynamically create and manage one or more virtual networks usable for communication between virtual network endpoints. In one example, virtual routers implement each virtual network using an overlay network, which provides the capability to decouple an endpoint's virtual address from a physical address (e.g., IP address) of the server on which the endpoint is executing. Each virtual network may use its own addressing and security scheme and may be viewed as orthogonal from the physical network and its addressing scheme. Various techniques may be used to transport packets within and across virtual networks over the physical network. At least some functions of the virtual router may be performed as one of services **233** or fabric service **235**. In the example of FIG. **2**, virtual router **220** executes within hypervisor **214** that uses physical function **221** for I/O, but virtual router **220** may execute within a hypervisor, a host operating system, a host application, one of virtual machines **228**, and/or processing unit **25** of NIC **230**.

[0058] In general, each virtual machine **228** may be assigned a virtual address for use within a corresponding virtual network, where each of the virtual networks may be associated with a different virtual subnet provided by virtual router **220**. A virtual machine **228** may be assigned its own virtual layer three (L3) IP address, for example, for sending and receiving communications but may be unaware of an IP address of the computing device **200** on which the virtual machine is executing. In this way, a “virtual address” is an address for an application that differs from the logical address for the underlying, physical computer system, e.g., computing device **200**.

[0059] In one implementation, computing device **200** includes a virtual network (VN) agent (not shown) that controls the overlay of virtual networks for computing device **200** and that coordinates the routing of data packets within computing device **200**. In general, a VN agent communicates with a virtual network controller for the multiple virtual networks, which generates commands to control routing of packets. A VN agent may operate as a proxy for control plane messages between virtual machines **228** and virtual network controller, such as controller **24** (FIG. **1**). For example, a virtual machine may request to send a message using its virtual address via the VN agent, and VN agent may in turn send the message and request that a response to the message be received for the virtual address of the virtual machine that originated the first message. In some cases, a virtual machine **228** may invoke a procedure or function call presented by an application programming interface of VN agent, and the VN agent may handle encapsulation of the message as well, including addressing.

[0060] In one example, network packets, e.g., layer three (L3) IP packets or layer two (L2) Ethernet packets generated or consumed by the instances of applications executed by virtual machine **228**

within the virtual network domain may be encapsulated in another packet (e.g., another IP or Ethernet packet) that is transported by the physical network. The packet transported in a virtual network may be referred to herein as an “inner packet” while the physical network packet may be referred to herein as an “outer packet” or a “tunnel packet.” Encapsulation and/or de-capsulation of virtual network packets within physical network packets may be performed by virtual router **220**. This functionality is referred to herein as tunneling and may be used to create one or more overlay networks. Besides IPinIP, other example tunneling protocols that may be used include IP over Generic Route Encapsulation (GRE), Virtual Extensible Local Area Network (VXLAN), Multiprotocol Label Switching (MPLS) over GRE (MPLSOGRE), MPLS over User Datagram Protocol (UDP)(MPLSoUDP), etc.

[0061] As noted above, a virtual network controller may provide a logically centralized controller for facilitating operation of one or more virtual networks. The virtual network controller may, for example, maintain a routing information base, e.g., one or more routing tables that store routing information for the physical network as well as one or more overlay networks. Virtual router **220** of hypervisor **214** implements a network forwarding table (NFT) **222A-222N** for N virtual networks for which virtual router **220** operates as a tunnel endpoint. In general, each NFT **222** stores forwarding information for the corresponding virtual network and identifies where data packets are to be forwarded and whether the packets are to be encapsulated in a tunneling protocol, such as with a tunnel header that may include one or more headers for different layers of the virtual network protocol stack. Each of NFTs **222** may be an NFT for a different routing instance (not shown) implemented by virtual router **220**.

[0062] In accordance with techniques of this disclosure, edge services controller **28** (FIG. **1**) uses processing unit **25** of NIC **230** to augment the processing and networking functionality of computing device **200**. Processing unit **25** includes processing circuitry **231** to execute services orchestrated by edge services controller **28**. Processing circuitry **231** may represent any combination of processing cores, ASICs, FPGAs, or other integrated circuits and programmable hardware. In an example, processing circuitry may include a System-on-Chip (SoC) having, e.g., one or more cores, a network interface for high-speed packet processing, one or more acceleration engines for specialized functions (e.g., security/cryptography, machine learning, storage), programmable logic, integrated circuits, and so forth. Such SoCs may be referred to as data processing units (DPUs). DPUs may be examples of processing unit **25**.

[0063] In the example NIC **230**, processing unit **25** executes an operating system kernel **237** and a user space **241** for services. Kernel **237** may be a Linux kernel, a Unix or BSD kernel, a real-time OS kernel, or other kernel for managing hardware resources of processing unit **25** and managing user space **241**.

[0064] Services **233** may include network, security, storage, data processing, co-processing, machine learning or other services. Services **233** and edge services platform (ESP) agent **236** include executable instructions. Processing unit **25** may execute instructions of services **233** and edge services controller (ESC) agent **236** as processes and/or within virtual execution elements such as containers or virtual machines. As described elsewhere in this disclosure, services **233** may augment the processing power of the host processors (e.g., microprocessor **210**), e.g., by enabling computing device **200** to offload packet processing, security, or other operations that would otherwise be executed by the host processors. Network services of services **233** may include security services (e.g., firewall), policy enforcement, proxy, load balancing, or other L4-L7 services.

[0065] Processing unit **25** executes ESC agent **236** to exchange data with edge services controller **28** (FIG. **1**) for the edge services platform. While shown in the example of FIG. **2** as being in user space **241**, in other examples, ESC agent **236** is a kernel module of kernel **237**. As an example, ESC agent **236** may collect and send telemetry data to the edge services controller **28** (FIG. **1**). The telemetry data may be generated by services **233** (FIG. **2**) and may describe traffic in the network,

availability of computing device **200** or network resources, resource availability of resources of processing unit **25** (such as memory or core utilization), or other information. As another example, ESC agent **236** may receive, from the edge services controller, service code to execute any of services **233**, service configuration to configure any of services **233**, packets or other data for injection into the network.

[0066] Edge services controller **28** (FIG. **1**) manages the operations of processing unit **25** by, e.g., orchestrating and configuring services **233** (FIG. **2**) that are executed by processing unit **25**, deploying services **233**; adding, deleting and replacing NICs within the edge services platform, monitoring of services **233** and other resources on NIC **230**, and managing connectivity between various services **233** running on NIC **230**. The edge services controller **28** may push one or more services **233** to NICs of one or more servers, based on based on predicted future application/service resource utilizations and server resource utilizations.

[0067] Example resources on NIC **230** include memory **227** and processing circuitry **231**. Edge services controller **28** may provide topology information via ESC agent **236**. Edge services controller **28** may provide flow information and/or forwarding information via ESC agent **236**. The flow information describes, and is usable for identifying, packet flows. The forwarding information is usable for mapping packets received by NIC **230** to an output port of NIC **230**.

[0068] FIG. **3** is a conceptual diagram illustrating a data center **300** with servers that each include a network interface card having a separate processing unit, controlled by an edge services platform, according to techniques of this disclosure. Racks of compute nodes **307A-307N** (collectively, “racks of compute nodes **307**”) may correspond to servers **12** of FIG. **1**, and switches **308A-308N** (collectively, “switches **308**”) may correspond to the switches of switch fabric **14** of FIG. **1**. An agent **302** of orchestrator **304** represents software executed by the processing unit (illustrated in FIG. **3** as a data processing unit or DPU) and receives configuration information for the processing unit and sends telemetry and other information for the NIC that includes the processing unit to orchestrator **304**. A gent **302** may represent an example of ESC agent **236**. Network services **312**, L4-L7 services **314**, telemetry service **316**, and Linux and software development kit (SDK) services **318** may represent examples of services **233**. Orchestrator **304** may represent an example of edge services controller **28** of FIG. **1**.

[0069] Network automation platform **306** connects to and manages network devices and orchestrator **304**, by which network automation platform **306** can utilize the edge services platform. Network automation platform **306** may, for example, deploy network device configurations, manage the network, extract telemetry, and analyze and provide indications of the network status.

[0070] FIG. **4** is a block diagram illustrating an example computing device that uses a network interface card having a separate processing unit, to perform services managed by an edge services platform according to techniques described herein. Although virtual machines are shown in this example, other instances of computing device **400** may also or alternatively run containers, native processes, or other endpoints for packet flows. Different types of vSwitches may be used, such as Open vSwitch or a virtual router (e.g., Contrail). Other types of interfaces between endpoints and NIC are also contemplated, such as tap interfaces, veth pair interfaces, etc.

[0071] FIG. **5** is a block diagram illustrating an example data center **540** having an edge services controller **525** for scheduling one or more services to one or more servers, according to techniques of this disclosure. Edge services controller **525** may represent an example instance of edge services controller **28** of FIG. **1**. In the illustrated example, single or multiple instances of services may be running on DPUs of any of three servers, which are represented by and include a corresponding first DPU **511**, a second DPU **512**, and a third DPU **513**. For example, a first instance of a first service **501** and a second instance of a second service **502** may be running on the first DPU **511**. Likewise, a first instance of a third service **503**, a second instance of the first service **504**, and a first instance of the second service **505** may be running on the second DPU **512**. Similarly, a first instance of a fourth service **506** and a second instance of the third service **507** may be running on

the third DPU **513**. One or more processors (“processor **527**”) of edge services controller **525** may be configured to implement a service/application scheduler **523** for scheduling one or more services on at least one of the first, second, or third DPUs **511**, **512** and **513**. The edge services platform controller **525** may also include a scheduler configurator **522** that receives and tracks any server taints and/or affinity configurations. Server taints can be used to mark a server that is in an unusable, unstable, or security-compromised state, so that the application scheduler **523** can avoid scheduling a service to the marked server. Affinity configuration, which may also be referred to as DPU pinning or “cache affinity”, enables a binding of a service, process or thread to a specific DPU, so that the service, process or thread will execute only on the designated DPU rather than on any DPU. Processor **527** can be operatively coupled to a memory **529** configured for storing a service profile database **524** and a metric collector database **521**. The service profile database **524** associates each of a plurality of services/applications with a corresponding resource utilization level.

[0072] Consider an illustrative scenario where a resource utilization of the first service is high, and the utilization reaches a peak in a periodic and/or predictable pattern. When the peak is reached, the first service may consume almost all available resources, and/or may compete with other services running on the same DPU, such as the first DPU **511**. These factors may cause performance issues with respect to the first service, and/or any other service running on the first DPU **511**, due to insufficient resources on the first DPU **511**. For example, assume that the service/application scheduler **523** is configured for scheduling the second instance of the second service **502** using a traditional scheduling approach based on current resource availability. As the first instance of the second service **505** is already running on the second DPU **512**, the service/application scheduler **523** is able to select from among the first DPU **511** and the third DPU **513** as target servers for scheduling. The service/application scheduler **523** may obtain resource availability information for the first DPU **511** and the third DPU **513** from the metric collector/database **521**. The service/application scheduler **523** may also check a requirements profile for the second service from the service profile database **524** to see if the first DPU **511** and/or the third DPU **513** have required and/or sufficient resources to run the second service.

[0073] Pursuant to a traditional scheduling approach, the service/application scheduler **523** lacks knowledge of resource utilization patterns for services and servers. For example, the service/application scheduler **523** may determine that the first DPU **511** and the third DPU **513** both have sufficient resources for running the second service, but the first DPU **511** has more resources than the third DPU **513**. Thus, the service/application scheduler **523** may simply schedule the second instance of the second service onto the first DPU **511**, in response to the first DPU **511** having more available resources than the third DPU **513**. However, if the first DPU **511** is selected for running the first service, a relatively large amount of resources may be consumed on a periodic basis, causing other services to starve for resources. Such a scheduling approach can cause performance issues for the first instance of the first service **501** and the second instance of the second service **502**, which are both running on the first DPU **511**. In the present example, the second service **502** running on the first DPU **511** is shown in dashed lines to indicate that this scheduling may be suboptimal and/or inefficient.

[0074] Performance issues may arise when resource utilization of the first instance of the first service **501** reaches a peak value as per its resource utilization pattern. For example, if the second instance of the second service **502** is deployed to the first DPU **511**, the second instance of the second service **502** may be forced to compute for resources of the first DPU **511** with the first instance of first service **501** and may be unable to obtain sufficient resources of the first DPU **511** to operate effectively. This scenario represents an existing or traditional approach to scheduling.

[0075] The edge services controller **525** may acquire knowledge associated with future resource requirements of services and future resource availabilities of servers, such as the first, second, and third DPUs **511**, **512** and **513**. Based on this acquired knowledge, the service/application scheduler

523 may schedule the second instance of second service **502** to the third DPU **513** instead of the first DPU **511**. For example, the second instance of second service **502** can be scheduled to the third DPU **513** where resource consumption of already running services may not be high, such that the third DPU **513** is predicted to have sufficient future resources to accommodate the second instance of second service **502**.

[0076] FIG. **6** is a data flow diagram illustrating an example method for scheduling one or more services. The procedure may commence at step **712** where a user/auto scaler **702** sends a request to the service/application scheduler **523** requesting a deployment and/or an auto-scaling of a service. At step **714**, the service/application scheduler **523** may obtain a service profile from the service profile database **524**. The service profile associates a specified service/application with a corresponding resource utilization level.

[0077] The service/application scheduler **523** may obtain server taints and/or an affinity configuration by user from the scheduler configurator **522** at step **716**. As mentioned previously, the server taint can be used to mark a server that is in an unusable, unstable, or security-compromised state, so that the application scheduler **523** can avoid scheduling a service to the marked server. Affinity configuration enables a binding of a service, process or thread to a specific DPU, so that the service, process or thread will execute only on the designated DPU rather than on any DPU. The service/application scheduler **523** may then select a server based on current resource availability, the server taints, and/or the affinity configuration at step **718**. At step **720**, pursuant to a traditional or conventional approach, the service/application scheduler **523** may discard a correct or optimum server for scheduling while selecting a server (i.e., a DPU such as the first, second, or third DPUs **511**, **512**, **513** of FIG. **5**) based solely on current resource availability. Accordingly, at step **722** (FIG. **7**), the service/application scheduler **523** schedules the service to a wrong (i.e., a less optimal or less efficient) server due to the scheduling being based solely on current resource availability.

[0078] FIG. **7** is a block diagram illustrating an example system for scheduling one or more services to one or more servers, according to techniques of this disclosure. An edge services controller **725** implements a self-learning scheduler. The self-learning scheduler may use a learning engine **705** to learn about resource utilization requirements of services, and resource utilization patterns of servers, to schedule services to appropriate servers to further optimize service scheduling. By analyzing historical usage data of servers and services/applications obtained from the metric collector/database **521**, the learning engine **705** can train a machine learning model to predict resource utilization values for servers and services/applications at any given timestamp.

[0079] Based on the predicted resource utilization patterns for each of a plurality of servers, a server weights predictor **703** of the learning engine **705** may calculate a server weight for each of the plurality of servers to thereby provide a set of server weights (Sw). Each respective server weight in the set of server weights (Sw) can be indicative of a resource availability on a corresponding server for scheduling. In one embodiment, high values of Sw for a given server indicate that more resources are available on the server, whereas low values of Sw indicate that less resources are available on the server. Based on the predicted resource utilization values of services/applications, an application weights predictor **701** of the learning engine **705** may calculate an application weight for each of one or more services/applications, to thereby provide a set of application weights (Aw). The set of application weights (Aw) may indicate the resource utilization requirements of the one or more services/applications. In one embodiment, high values of Aw indicate that the service requires more resources on the server, whereas low values of Aw indicate that the service requires less resources on the server. The learning engine **705** may instruct the service/application scheduler **523** to schedule a service based on the calculated set of server weights (Sw) and the calculated set of application weights (Aw). In one embodiment, a service with a high Aw is scheduled to a server with a high Sw. Accordingly, a service requiring high resources can be scheduled to a server having more resources.

[0080] Pursuant to an illustrative example, assume that a backup service is to be scheduled by a user every two hours throughout a 24-hour day. The resource utilization of the backup service may reach a peak every two hours. The machine learning model of the learning engine **705** analyzes the resource utilization metrics of the backup service as received from the metric collector/database **521**, as well as the resource utilization metrics of the servers on which the backup service runs, to predict one or more timestamps at which the backup service would consume a peak and/or maximum amount of resources, thereby causing servers to have less resources available for other co-located services.

[0081] Resource usage metrics of individual services and servers can be exported from the metric collector/database **521** as a time series database (e.g., Prometheus). The collected resource usage metric data can be used to train the machine learning model of the learning engine **705** to predict the resource usage metrics at future timestamps. The trained machine learning model can be deployed to formulate a solution using the learning engine **705**. The learning engine **705** can continue to analyze live resource usage data from the metric collector/database **521** and train the already-trained machine learning model to improve overall accuracy of the predictions.

[0082] The service/application scheduler **523**, informed by one or more predictions from learning engine **705**, may schedule the second instance of the second service **502** to the third DPU **513** rather than the first DPU **511**. This is because the learning engine **705** possesses knowledge encompassing future resource requirements of services as well as resource availabilities of servers. Thus, the learning engine **705** can schedule the second instance of S2 onto the third DPU **513** where resource consumption of already running services not going to be high and the server (third DPU **513**) has the required resources to accommodate the second instance of the second service **502**, without causing any performance issues.

[0083] Some implementations described herein use learning engine **705** to generate a machine-learning based solution configured to predict future requirements of services, and also configured to predict future resource availabilities of servers. For example, some implementations described herein may train a model using a machine learning technique. The model may be trained based on observed operational information (e.g., telemetry data and/or the like) for a set of DPUs **511**, **512**, **513** and based on flow information for traffic flows processed by the set of DPUs while running specific applications and/or services. The model may output predicted performance information for the set of DPUs **511**, **512**, **513** based on input information identifying traffic flows and/or operational information.

[0084] Furthermore, according to some implementations described herein, the learning engine **705** may update the model using the machine learning techniques and based on observations regarding efficacy of any of the set of DPUs **511**, **512**, **513** executing services and/or applications. In this way, the model may adapt to changing DPU conditions and topology (e.g., in real time as the network conditions and/or the topology change). Thus, throughput, reliability, and conformance with SLAs is improved. Further, some implementations described herein may use a rigorous, well-defined approach to service/application scheduling, which may reduce uncertainty, subjectivity, and inefficiency that may be introduced by a human actor attempting to define a scheduling policy based on empirical observations regarding network and DPU performance.

[0085] A Iso, some implementations described herein may identify a best or optimum DPU for executing an application/service at a given timestamp. Since the best or optimum DPU may iteratively change based on DPU load and DPU behavior/faults, the machine learning component of implementations described herein may regularly re-predict the best or optimum DPU to optimize application/service execution at any given time. This reprogramming may be based on dynamic prediction of DPU load, DPU dropouts, and/or DPU delays. Thus, implementations described herein may improve adaptability and versatility of DPU scheduling for applications/services in comparison to a rigidly defined scheduling protocol.

[0086] Furthermore, by using machine learning, implementations described herein may predict

DPU delays or dropouts, or reduced capacity on network devices, and may perform pre-emptive scheduling updates to avoid inefficiencies or dropouts due to DPU degradation. Thus, forward-looking maintenance and scheduling is provided, which further improves reliability and performance.

[0087] The model trained by learning engine **705** and used for inference/prediction, as described in this disclosure, may be a machine learning (ML) model. Learning engine **705** may be able to train various types of ML models. For instance, in some examples, learning engine **705** is configured to train baseline ML models. A baseline ML model may be a type of ML model other than a deep learning ML models and statistical ML models. Baseline ML models may be able to generate predictions based on limited amounts of data. For example, a baseline ML model may be able to generate a prediction based on less than 1 hour of data (e.g., for hourly predictions). Example types of baseline M L models may include an Exponential Weighted Moving Average (EWMA) model, a Hidden Markov model, and so on.

[0088] In some examples, learning engine **705** is configured to train statistical ML models. Example types of statistical models include a Holt-Winters model, an autoregressive integrated moving average (ARIMA) model, a seasonal ARIMA model, a vector autoregression (VAR) model, a Facebook PROPHET model, and so on. In some examples, statistical ML models may have greater utility than basic ML models when there is more data available to use to make predictions. For instance, a statistical ML model that is used to generate hourly predictions may be usable when more than 24 hours of data is available.

[0089] In some examples, learning engine **705** is configured to train deep learning ML models. Deep learning ML models may require more data than basic ML models or statistical ML models but may be able to provide more sophisticated types of predictions. Example types of deep learning ML models may include Long Short-Term Memory (LSTM) models, bi-directional LSTM models, recurrent neural networks, or other types of neural networks that include multiple layers. In other examples, learning engine **705** may use neural network models other than deep learning ML models

[0090] The ML models may be grouped as regression-based ML models, classification-based ML models, and unsupervised learning models. There may be baseline, statistical, and deep learning MLs for each of these groups. In some examples, for regression-based ML models, learning engine **705** may use a Hodrick-Prescott filter to perform an initial level of ML model selection.

Specifically, the Hodrick-Prescott filter breaks time-series data (y_t) into a trend component and a cyclical component c_t : $y_t = \text{trend}_t + c_t$ (trend)+ c_t (cyclical). The time-series data is the data that the ML models use to generate the predictions. By breaking the time-series data into a trend component and a cyclical component, learning engine **705** may be able to determine whether the time-series data has more of a cyclic nature or more of a trend nature and use an appropriate ML model based on the determination. For example, the EWMA model and Holts-Winter model perform better on time-series data that has a cyclic nature. An ARIMA model, a VAR model, etc., may perform better on time-series data that has a trend nature.

[0091] By performing this initial level of ML model selection, learning engine **705** may be able to avoid training every regression-based ML model, thereby potentially saving time and computational resources. In some examples, learning engine **705** may filter the regression-based M L models based on how much data is available. For instance, if there is less than a threshold amount of time's worth of available training data (e.g., 24-48 hours), learning engine **705** may train only regression-based baseline ML models. Otherwise, if there is more than the threshold amount of time's worth of available data, learning engine **705** may additionally or alternatively train other types of regression-based ML models, such as statistical models or low capacity deep learning ML models.

[0092] Example types of regression-based baseline ML models may include a hidden Markov model and season trend decomposition approaches. Example types of regression-based statistical

ML models may include Error-Trend-Seasonality (ETS) models (including exponential smoothing models, trend method models, and ETS decomposition), EWMA models (including simple moving averages and EWMA), Holt Winters models, ARIMA models, SARIMA models, vector autoregression models, seasonal trend autoregression (STAR) models, and Facebook PROPHET models. Example types of regression-based deep learning ML models may include LSTM architectures (including single-layer LSTM s, depth LSTM s, bi-directional LSTM s), RNNs, and gated recurrent units (GRUs). Example types of classification-based baseline M L models may include logistic regression models and K-nearest neighbor models. Example types of classification-based statistical ML models may include support vector machines and boosting ensemble algorithms (e.g. XGBoost). Example types of classification-based deep learning ML models may include LSTM architectures, RNN architectures, GRU architectures, and artificial neural network architectures. Example types of unsupervised ML models may include K-means clustering models, Gaussian clustering models, and density-based spatial clustering.

[0093] Prediction of the Set of Application Weights (Aw): The application weight can be an attribute of every service scheduled to run on any of the servers. It may be an indicator of resource usage requirements of a service at any given timestamp. When a service instance is going to consume more resources at a next hour, a next minute, or a next second, a higher value can be assigned to the application weight (A w) compared to other less resource consuming services. Thus, when a service instance has a higher application weight (Aw) value, this may be indicative that the service instance would need more resources at the given timestamp. The application weight (Aw) can be predicted by the trained machine learning model using resource usage telemetry data of a service relative to a starting time of the service/application, as received from the metric collector/database 521. The machine learning model can predict the resource usage values for the given timestamp, and the set of application weights is calculated based on predicted resource usage metric values including one or more of: CPU Usage, Network Usage, DPU Usage, or another metric. In some embodiments, any one or more of the following resource usage telemetry metrics can be used to predict the set of application weights (A w): [0094] 1) CPU Usage; [0095] 2) Memory Usage; and/or [0096] 3) Network Usage.

[0097] The historical data values of any of these metrics can be used to predict the metric values at different future timestamp values. Using the historical data of above metrics, a machine learning model is trained to predict metric values at future timestamps. In some examples, the machine learning model is implemented using Vector Auto Regression.

[0098] An example mathematical model for a metric value predictor machine learning model is: [00001] $MV(t) = a + W1 * MV(t - 1) + .Math. + Wp * MV(t - p) + e(t)$ [0099] where: [0100] MV is Metric Value at future time t [0101] a is constant [0102] Wl & Wp are co-efficient values [0103] p is number of past values of metric [0104] e(t) is error correction at time t

[0105] Table 1 shows a historical or past data of collected metrics and forecasted values at different future timestamps.

TABLE-US-00001

TABLE 1	Day (N-2)	Day (N-1)	Day N	CPU	DPU	Network	CPU	DPU	Network
CPU	DPU	Network	Time*	Usage	Usage	Usage	Usage	Usage	Usage
t + Dt1	3	2	4	2	4	1	3	4	2
t + Dt2	12	8	7	13	6	8	11	7	6
t + Dt3	7	5	6	5	7	4	6	5	5

*-> Time is relative to start time of service instance. Note: To simplify, this table only shows the learning process for some of the metrics. The same process can be applied to a different set of metrics in various examples.

[0106] Table 1 shows a resource usage pattern of a service. After obtaining forecasted values of resource usage metrics of a service/application at different timestamp values relative to start timestamp of service, a weight fraction can be calculated for each individual metric. For example, a 100% CPU usage metric may contribute 20% and a 100% DPU usage metric may contributes 40% to an application weight of the set of application weights (A w). The CPU usage metric value of 30% may contribute 0.06 to the application weight, and 40% DPU usage may contribute 0.2 to the application weight.

[0107] Application Weight for Metric (AWmx)=(Metric Value*Metric Weight Factor)/100, where Metric Value is predicted value of a metric, Metric Weight Factor is user adjusted or configured fraction between 0 and 1. One may take a mean of the predicted weights as a weight of the server: [00002]ApplicationWeight(Aw) = (.Math.ⁿ AWmx1 + AWmx2 + .Math. AWmxi)

[0108] Prediction of the set of Server Weights (Sw): The set of server weights (Sw) may comprise an attribute of every server managed by the ESP Controller 725 and/or an orchestrator. This may be an indicator of cumulative resource usage by all service instances running on the server at any given timestamp. When resource usage metrics of a server are increasing in a next hour, a next minute or a next second, a higher value can be assigned to a server weight in the set of server weights (Sw) compared to other servers whose resource usage metrics are lower. When a server has a higher server weight value, this may indicate that resource usage will be higher on that server at a predicted timestamp.

[0109] The set of server weights (Sw) can be predicted by the trained machine learning model using resource usage telemetry data of a service instance acquired from the metric collector/database 521. Any one or more of the following resource usage telemetry metrics can be used to predict the set of server weights (Sw): [0110] 1 CPU Usage: Measured as a percentage [0111] 2 Network Usage: Measured as a percentage [0112] 3 DPU Usage: Measured as a percentage [0113] 4 Jumbo packets [0114] 5 ECMP usage percentage [0115] 6 Elephant flows [0116] 7 Mice flows [0117] 8 Traffic leaving DC [0118] 9 Encrypted traffic % [0119] 10 Compressed traffic % [0120] 11 Number of firewall rules [0121] 12 Network Latency: as milliseconds [0122] 13 Network Packet Loss: as percentage [0123] 14 Network Throughput: as bps [0124] 15 Jitter: as milliseconds

[0125] As shown in FIG. 10, metric values for a present day can be forecast using metric values of a last two days. To simplify, the learning and prediction process for a few illustrative metrics is shown. This process can be applied to any additional metrics as well.

[0126] FIG. 10 shows an illustrative resource utilization pattern on a DPU (such as the first DPU 511). The metric data of last 2 days indicates that CPU, network and DPU usage and request rate values are high towards the end of the day, so a probability of the machine learning model predicting higher values for these metrics is high.

[0127] An example mathematical model for a metric value predictor machine learning model is: [00003] $MV(t) = a + W1 * MV(t - 1) + .Math. + Wp * MV(t - p) + e(t)$ [0128] Where: [0129] MV is Metric Value at future time t [0130] a is constant [0131] W1 & Wp are co-efficient values [0132] p is number of past values of metric [0133] e(t) is error correction at time t

[0134] After obtaining forecasted values of resource usage metrics of a server at different timestamp values, a weight fraction can be calculated for each individual metric. For example, assume that a 100% CPU usage metric contributes 20% and a 100% DPU usage metric contributes 40% to a server weight. The CPU usage metric value of 30% can contribute 0.06 to the server weight and 40% DPU Usage can contribute 0.2 to the server weight

[00004]ServerWeightforMetric(AWmx) = (MetricValue * MetricWeight Factor) / 100, [0135] where [0136] Metric Value is predicted value of a metric [0137] Metric Weight Factor is user adjusted or configured fraction between 0 and 1.

[0138] One may take a mean of predicted weights as a weight of a server:

[00005]ServerWeight(Aw) = (.Math.ⁿ AWmx1 + AWmx2 + .Math. AWmxi)

[0139] The edge services platform 725 can schedule a service based on the calculated set of server weights (Sw) and the calculated set of application weights (Aw). In one embodiment, the edge services platform 725 schedules a service with a high A w onto a server with a high Sw. Accordingly, a service requiring high resources can be deployed on a server having more resources. [0140] FIG. 11 shows a server and application weight matrix for multiple data processing units of servers at each of a plurality of different time stamps. A second instance 'S2-2' of a second service S2 can be scheduled to the third DPU 513 instead of the first DPU 511. That is because the third

DPU **513** has a highest server weight complementing a higher application weight A_w of the second service **S2-2** at the time of scheduling of **S2-2**.

[0141] FIG. **8** is a data flow diagram illustrating an example method for scheduling one or more services to one or more servers, according to techniques of this disclosure. At step **812**, the metric collector/database **521** receives server and application metrics from a server **810**. At step **814**, the learning engine **705** trains the machine learning model using one or more metrics obtained from the metric collector/database **521**. The learning engine **705** predicts server metric values and the set of server weights (S_w) at step **816**. The learning engine **705** predicts application metric values and computes a set of application weights (A_w) at step **818**. The learning engine **705** assigns predicted server weights of the set of server weights (S_w) to corresponding servers and predicted application weights of the set of application weights (A_w) to corresponding application instances at step **820**. At step **822**, the service/application scheduler **523** obtains at least one server weight of the set of server weights (S_w), and at least one application weight of the set of application weights (A_w), from the learning engine **705**. At step **824**, the service/application scheduler **523** selects a server and an application/service instance using the at least one server weight and the at least one application weight. An optional prior selection of a wrong server may be discarded, and an application/server instance is scheduled by the learning engine **705** to the correct server at step **826**.

[0142] FIG. **9** is a flowchart of an example method performed by an edge services platform controller **725** (FIG. **7**) according to techniques of this disclosure. At block **902** (FIG. **9**), a respective server weight is determined for each of a plurality of servers, based on a respective predicted resource utilization value for each of the plurality of servers. At block **904**, a respective application weight is determined for each of a plurality of services, based on a respective predicted resource utilization value for each of the plurality of services. At block **906**, a service of the plurality of services is scheduled on a server of the plurality of servers based on the respective server weight for the server and the respective application weight for the service.

[0143] Like reference characters denote like elements throughout the description and figures.

[0144] The techniques described herein may be implemented in hardware, software, firmware, or any combination thereof. Various features described as modules, units or components may be implemented together in an integrated logic device or separately as discrete but interoperable logic devices or other hardware devices. In some cases, various features of electronic circuitry may be implemented as one or more integrated circuit devices, such as an integrated circuit chip or chipset.

[0145] If implemented in hardware, this disclosure may be directed to an apparatus such as a processor or an integrated circuit device, such as an integrated circuit chip or chipset. Alternatively or additionally, if implemented in software or firmware, the techniques may be realized at least in part by a computer-readable data storage medium comprising instructions that, when executed, cause a processor to perform one or more of the methods described above. For example, the computer-readable data storage medium may store such instructions for execution by a processor.

[0146] A computer-readable medium may form part of a computer program product, which may include packaging materials. A computer-readable medium may comprise a computer data storage medium such as random access memory (RAM), read-only memory (ROM), non-volatile random access memory (NV RAM), electrically erasable programmable read-only memory (EEPROM), Flash memory, magnetic or optical data storage media, and the like. In some examples, an article of manufacture may comprise one or more computer-readable storage media.

[0147] In some examples, the computer-readable storage media may comprise non-transitory media. The term “non-transitory” may indicate that the storage medium is not embodied in a carrier wave or a propagated signal. In certain examples, a non-transitory storage medium may store data that can, over time, change (e.g., in RAM or cache). The code or instructions may be software and/or firmware executed by processing circuitry including one or more processors, such as one or more digital signal processors (DSPs), general purpose microprocessors, application-specific integrated circuits (ASICs), field-programmable gate arrays (FPGAs), or other equivalent

integrated or discrete logic circuitry. Accordingly, the term “processor,” as used herein may refer to any of the foregoing structure or any other structure suitable for implementation of the techniques described herein. In addition, in some aspects, functionality described in this disclosure may be provided within software modules or hardware modules.

Claims

1. One or more computing devices comprising: processing circuitry having access to memory storing executable instructions for a learning engine, the learning engine configured to: obtain first historical utilization data for a plurality of servers, the first historical utilization data comprising respective central processing unit (CPU) utilization at first times and respective data processing unit (DPU) utilization at second times, wherein the CPU utilization is for one or more CPUs, and wherein the DPU utilization is for one or more DPUs of one or more network interface cards (NICs); obtain second historical utilization data for a plurality of services, the second historical utilization comprising respective resource utilization requirements for the plurality of services at third times; and process the first historical utilization data and the second historical utilization data to train a machine learning model to predict: CPU utilization at a future time; DPU utilization at the future time; and a resource utilization requirement for a first service of the plurality of services at the future time.
2. The one or more computing devices of claim 1, wherein the processing circuitry is configured to output an instruction to schedule the first service at the future time based on a prediction of the trained machine learning model.
3. The one or more computing devices of claim 1, wherein the processing circuitry is configured to iteratively train the trained machine learning model with a current DPU utilization for one or more DPUs of the one or more NICs.
4. The one or more computing devices of claim 1, wherein the processing circuitry is configured to iteratively train the trained machine learning model with a current resource utilization requirement for the first service.
5. The one or more computing devices of claim 1, wherein the first service comprises one of a network, security, storage, data processing, co-processing, or machine learning service.
6. The one or more computing devices of claim 1, wherein the DPU utilization at the future time comprises an indication of availability of one or more of a processing core or a memory of a DPU of the one or more DPUs.
7. The one or more computing devices of claim 1, wherein the machine learning model comprises a vector autoregression machine learning model.
8. A method comprising: obtaining first historical utilization data for a plurality of servers, the first historical utilization data comprising respective central processing unit (CPU) utilization at first times and respective data processing unit (DPU) utilization at second times, wherein the CPU utilization is for one or more CPUs, and wherein the DPU utilization is for one or more DPUs of one or more network interface cards (NICs); obtaining second historical utilization data for a plurality of services, the second historical utilization comprising respective resource utilization requirements for the plurality of services at third times; and processing, by one or more computing devices, the first historical utilization data and the second historical utilization data to train a machine learning model to predict: CPU utilization at a future time; DPU utilization at the future time; and a resource utilization requirement for a first service of the plurality of services at the future time.
9. The method of claim 8, further comprising: outputting an instruction to schedule the first service at the future time based on a prediction of the trained machine learning model.
10. The method of claim 8, further comprising: iteratively training, by the one or more computing device, the trained machine learning model with a current DPU utilization for one or more DPUs of

the one or more NICs.

11. The method of claim 8, further comprising: iteratively training, by the one or more computing device, the trained machine learning model with a current resource utilization requirement for the first service.

12. The method of claim 8, wherein the first service comprises one of a network, security, storage, data processing, co-processing, or machine learning service.

13. The method of claim 8, wherein the DPU utilization at the future time comprises an indication of availability of one or more of a processing core or a memory of a DPU of the one or more DPUs.

14. The method of claim 8, wherein the machine learning model comprises a vector autoregression machine learning model.

15. Non-transitory computer-readable media comprising instructions that, when executed, cause processing circuitry to: obtain first historical utilization data for a plurality of servers, the first historical utilization data comprising respective central processing unit (CPU) utilization at first times and respective data processing unit (DPU) utilization at second times, wherein the CPU utilization is for one or more CPUs, and wherein the DPU utilization is for one or more DPUs of one or more network interface cards (NICs); obtain second historical utilization data for a plurality of services, the second historical utilization comprising respective resource utilization requirements for the plurality of services at third times; and process the first historical utilization data and the second historical utilization data to train a machine learning model to predict: CPU utilization at a future time; DPU utilization at the future time; and a resource utilization requirement for a first service of the plurality of services at the future time.

16. The non-transitory computer-readable media of claim 15, wherein the instructions cause the processing circuitry to: output an instruction to schedule the first service at the future time based on a prediction of the trained machine learning model.

17. The non-transitory computer-readable media of claim 15, wherein the instructions cause the processing circuitry to: iteratively train the trained machine learning model with a current DPU utilization for one or more DPUs of the one or more NICs.

18. The non-transitory computer-readable media of claim 15, wherein the instructions cause the processing circuitry to: iteratively train the trained machine learning model with a current resource utilization requirement for the first service.

19. The non-transitory computer-readable media of claim 15, wherein the first service comprises one of a network, security, storage, data processing, co-processing, or machine learning service.

20. The non-transitory computer-readable media of claim 15, wherein the DPU utilization at the future time comprises an indication of availability of one or more of a processing core or a memory of a DPU of the one or more DPUs.
